

Executive Briefing: What You Need to Know about Fast Data

Dean Wampler, Ph.D.
dean@lightbend.com
[@deanwampler](https://twitter.com/deanwampler)



Photographs © Dean Wampler, 2007-2018 (except company and product logos and other photos where noted). All rights reserved. All other content © Lightbend, 2013-2018. All rights reserved.
Photograph: Tower Bridge, London.
Last update: May 19, 2018

You can download this and my other talks from polyglotprogramming.com/talks.



Based on
this report

lightbend.com/fast-data-platform

O'REILLY®

Fast Data Architectures for Streaming Applications

Getting Answers Now from
Data Sets that Never End

Dean Wampler

Compliments of
Lightbend



What We'll Discuss

- 
- Why streaming? Why now?
 - How should I choose technologies?
 - How will this impact the rest of my organization?

What We'll Discuss



Why Streaming?

Why should we bother in the first place.

- 
- New opportunities that require streaming
 - Upgrading batch applications for competitive advantage

Why Streaming?

Let's discuss two classes of applications, opportunities, and business drivers. One is fairly obvious; some new things have low latency requirements, so you have to act on data as it arrives. Others have been successful with traditional batch approaches, but it's a competitive advantage to accelerate them through streaming.



Fast Data Use Cases

Predictive Analytics

Apply ML models to large volumes of device data to pre-empt failures / outages



Real-time Personalization

Real-time marketing based on behavior, location, inventory levels, product promotions, etc.



Real-time Financial Processes

Drive better business outcomes through real-time risk, fraud detection, compliance, audit, governance, etc.



IoT

Real-time consumer and industrial Device and Supply Chain management at scale

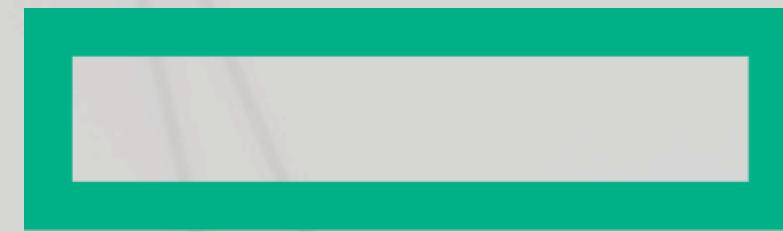


Legacy Modernization

Accelerate decision making processes and optimize infrastructure costs by moving from batch to streaming



More at: <https://www.lightbend.com/customers>

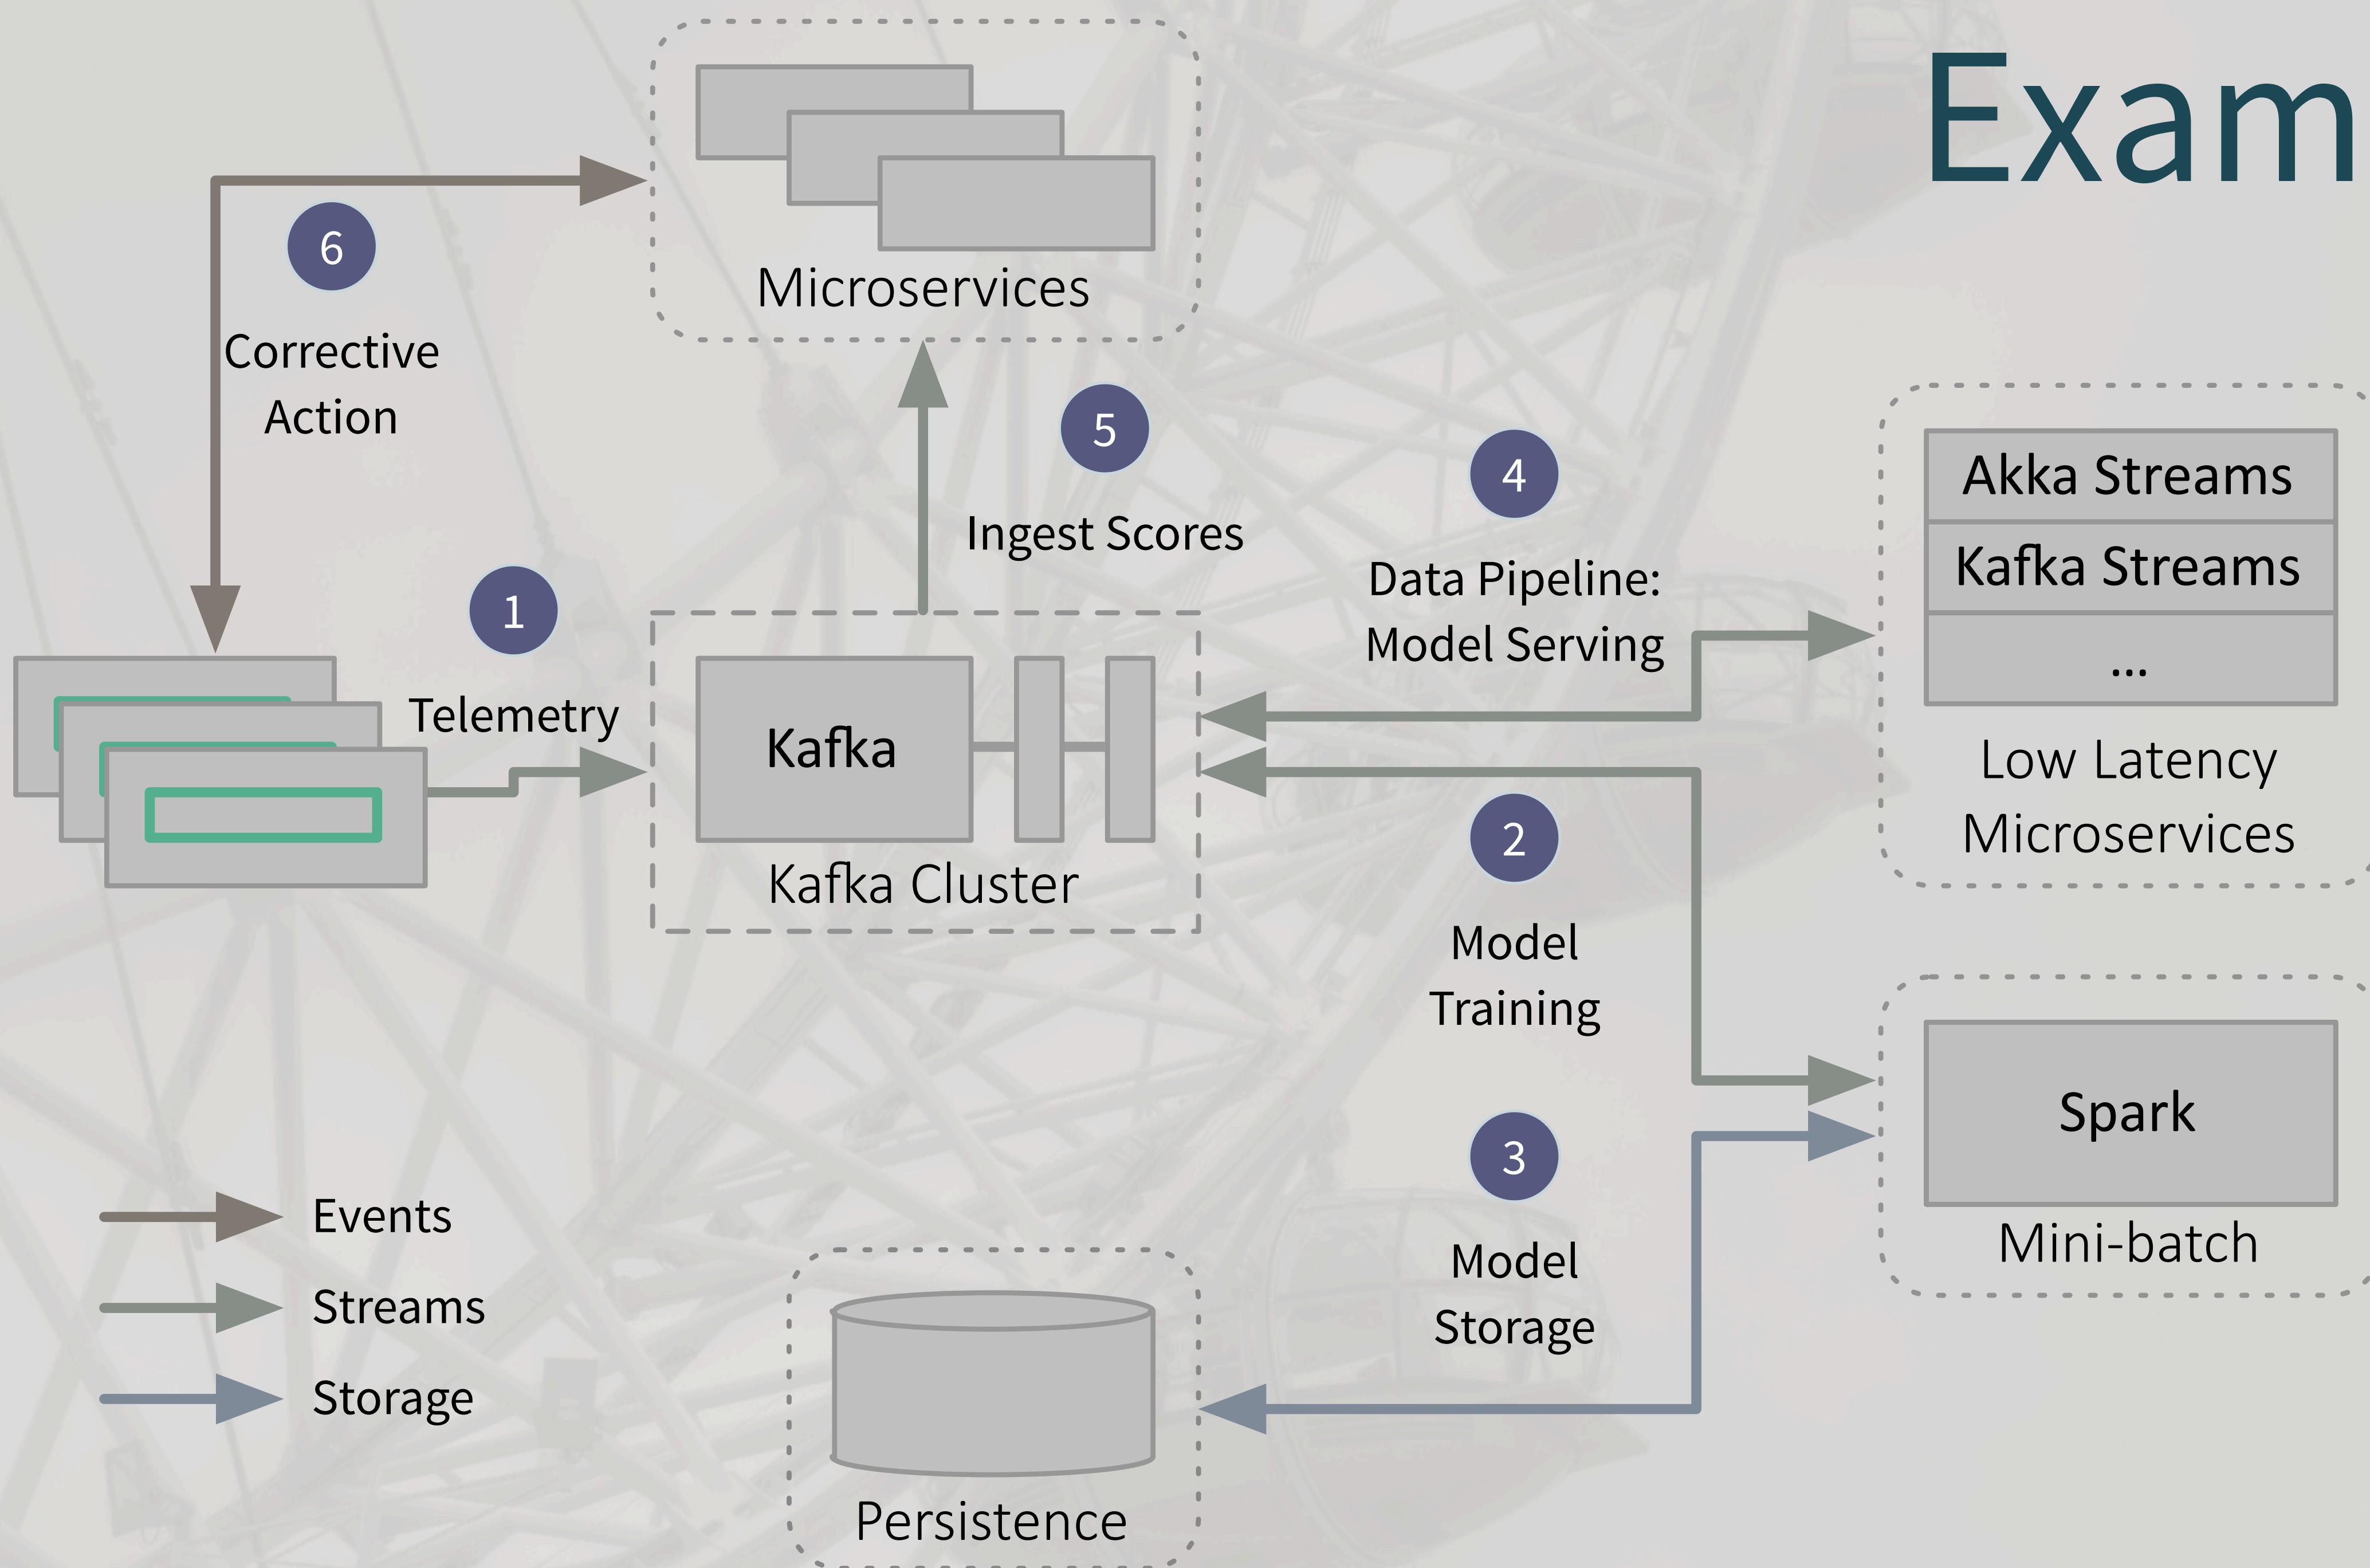


Predictive Analytics

Hewlett Packard Enterprise

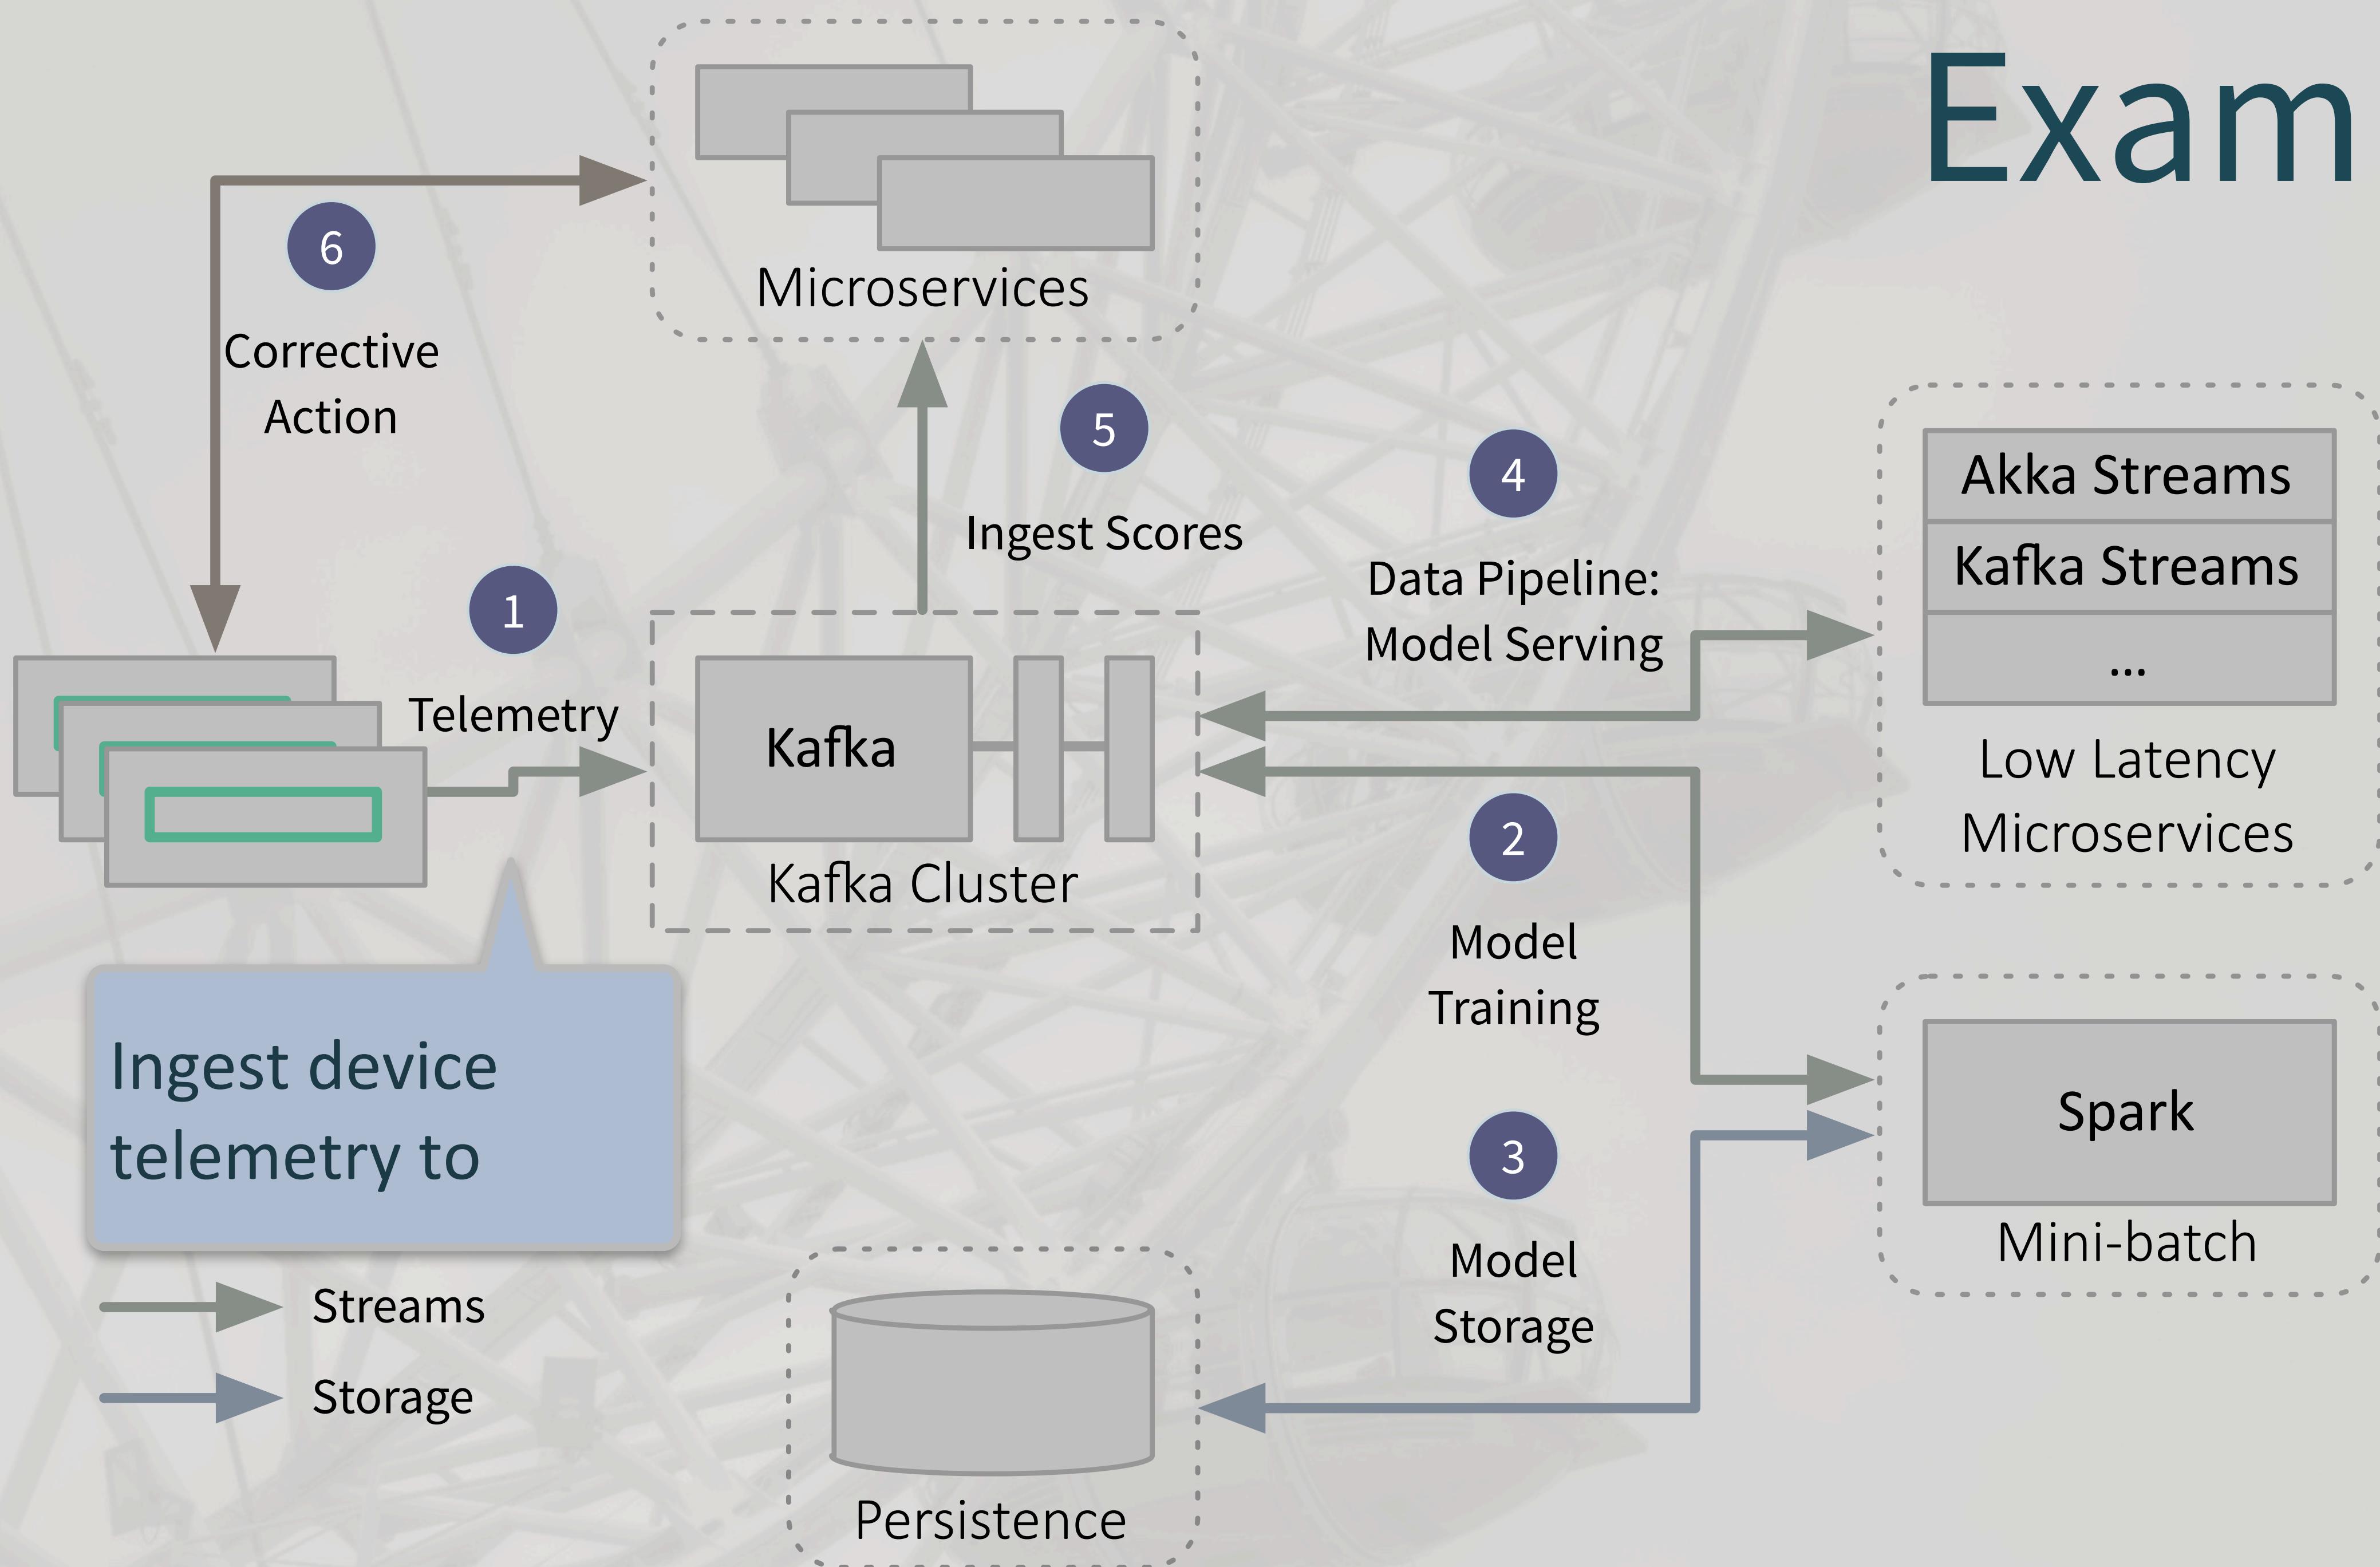
- ML models applied to device telemetry to detect anomalies
- Preemptive maintenance prevents user impacts from potential failures

Example



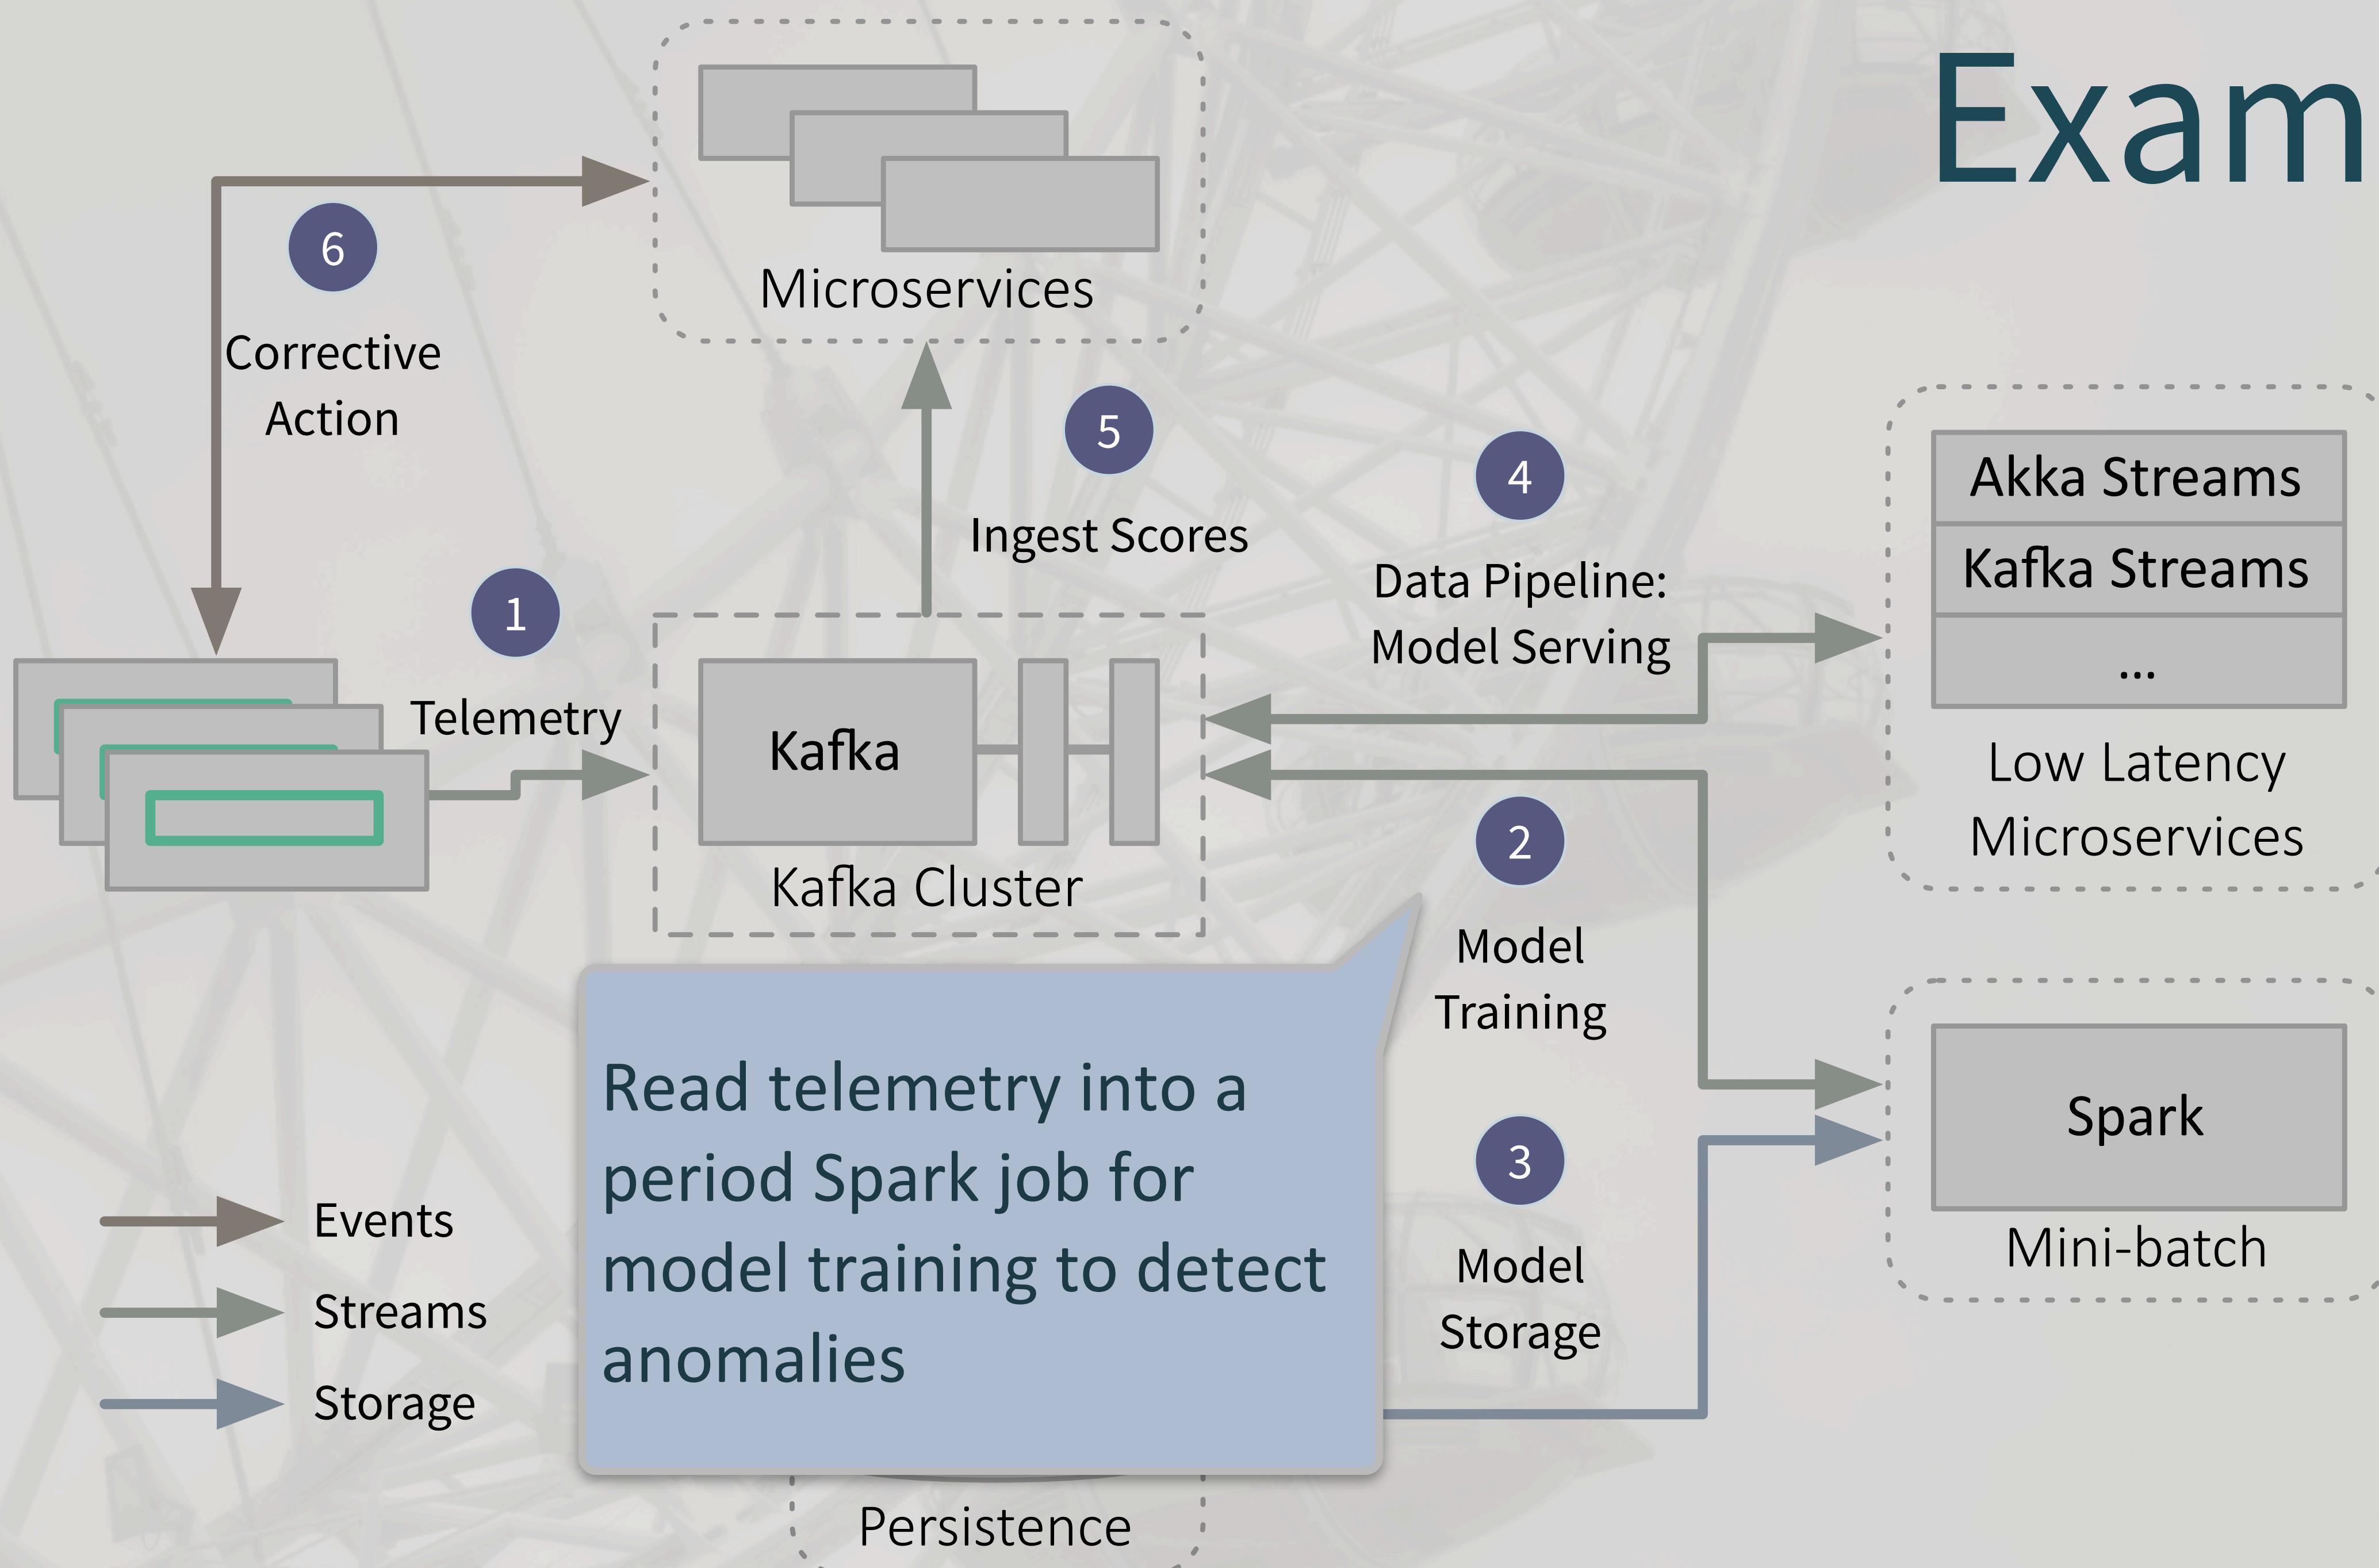
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



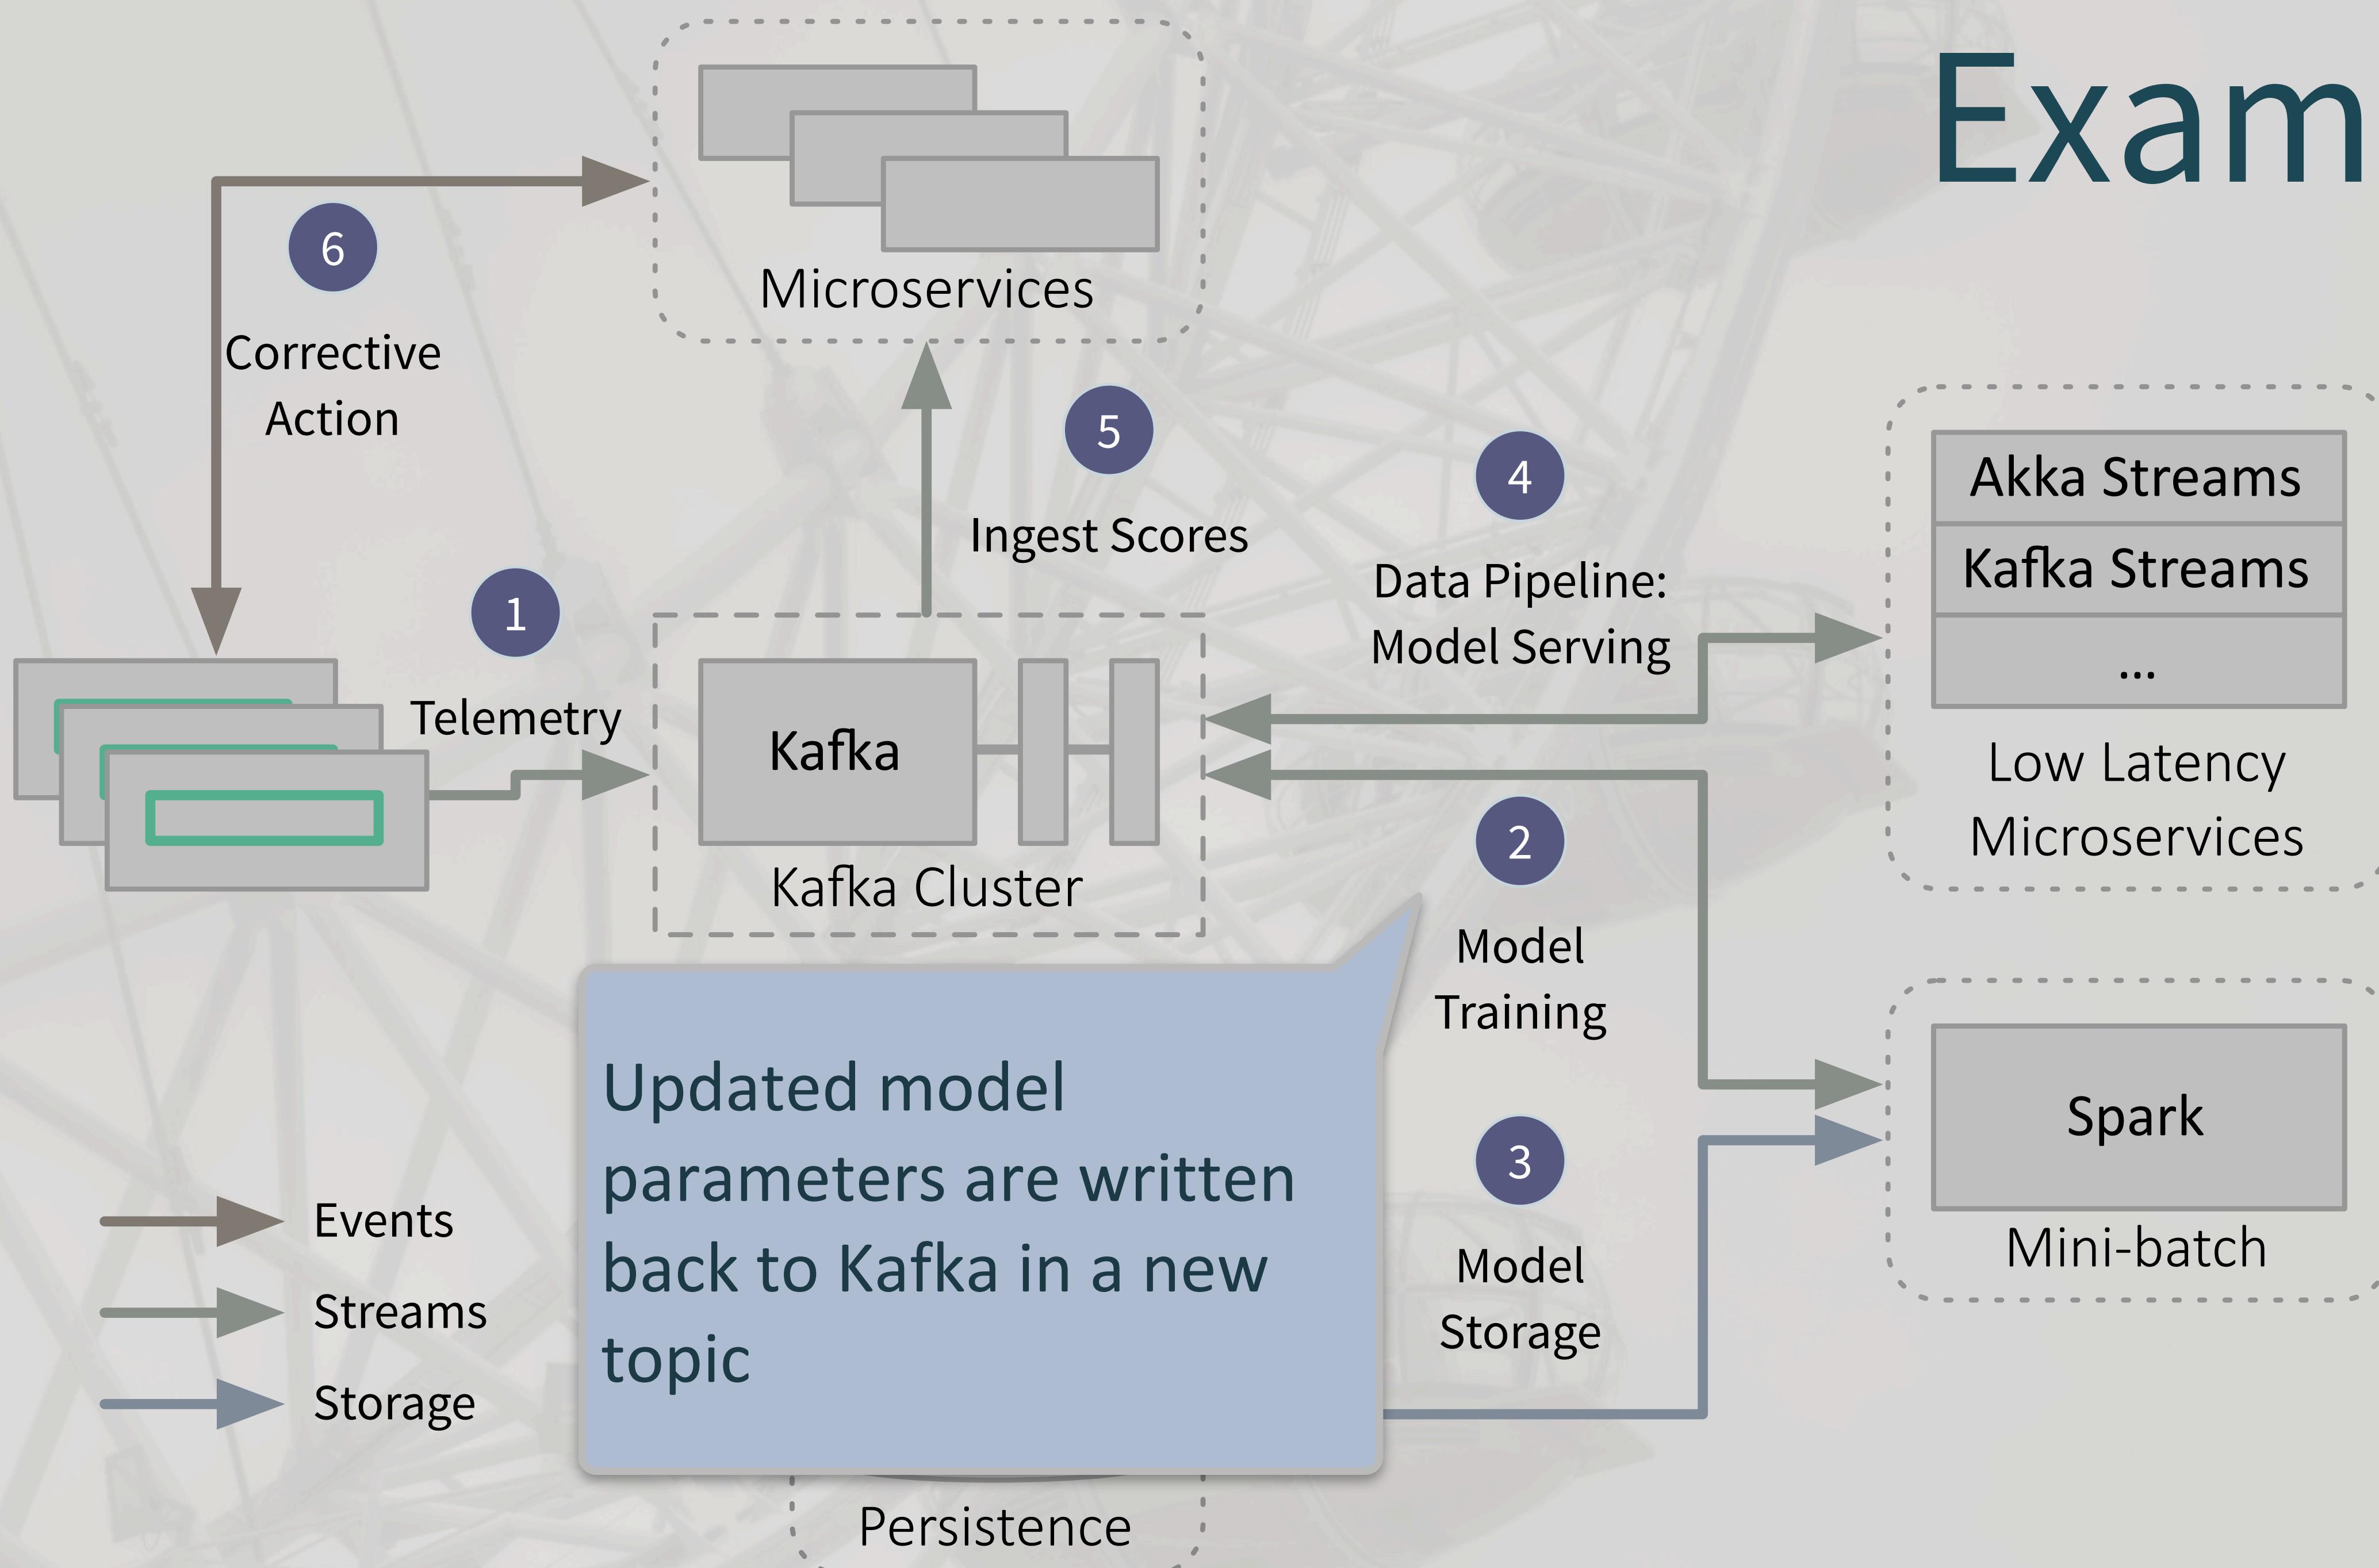
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



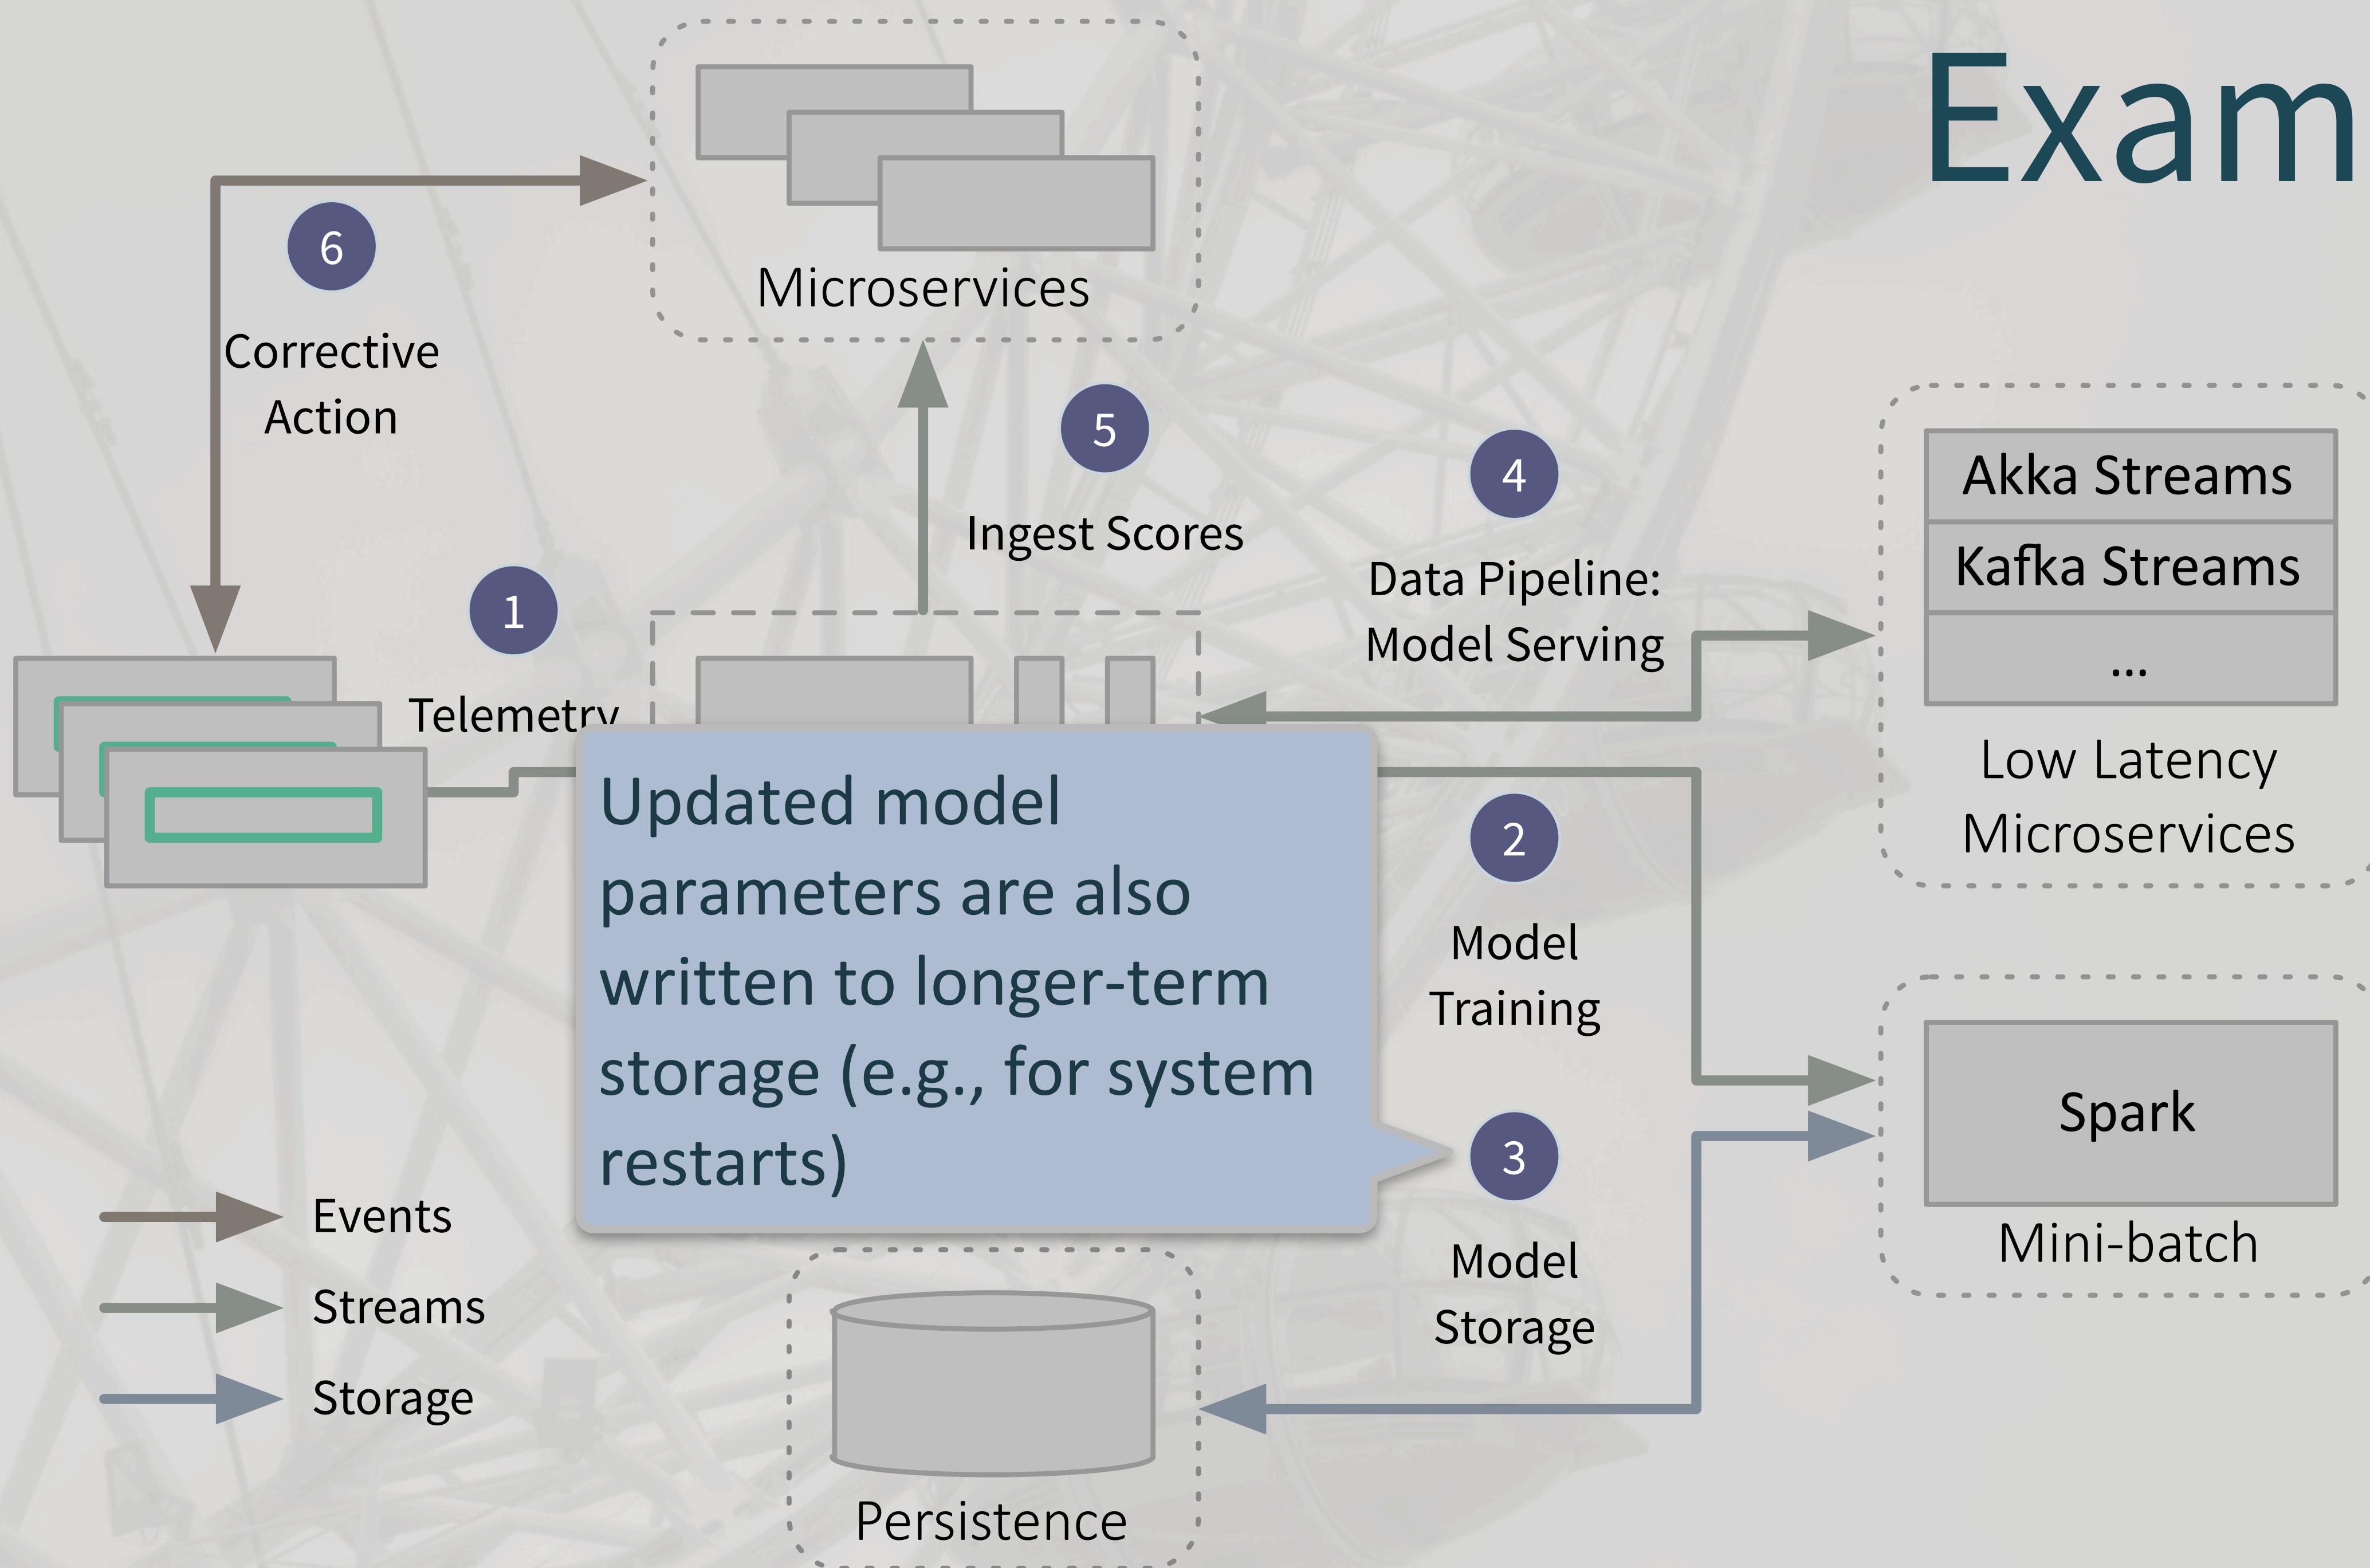
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



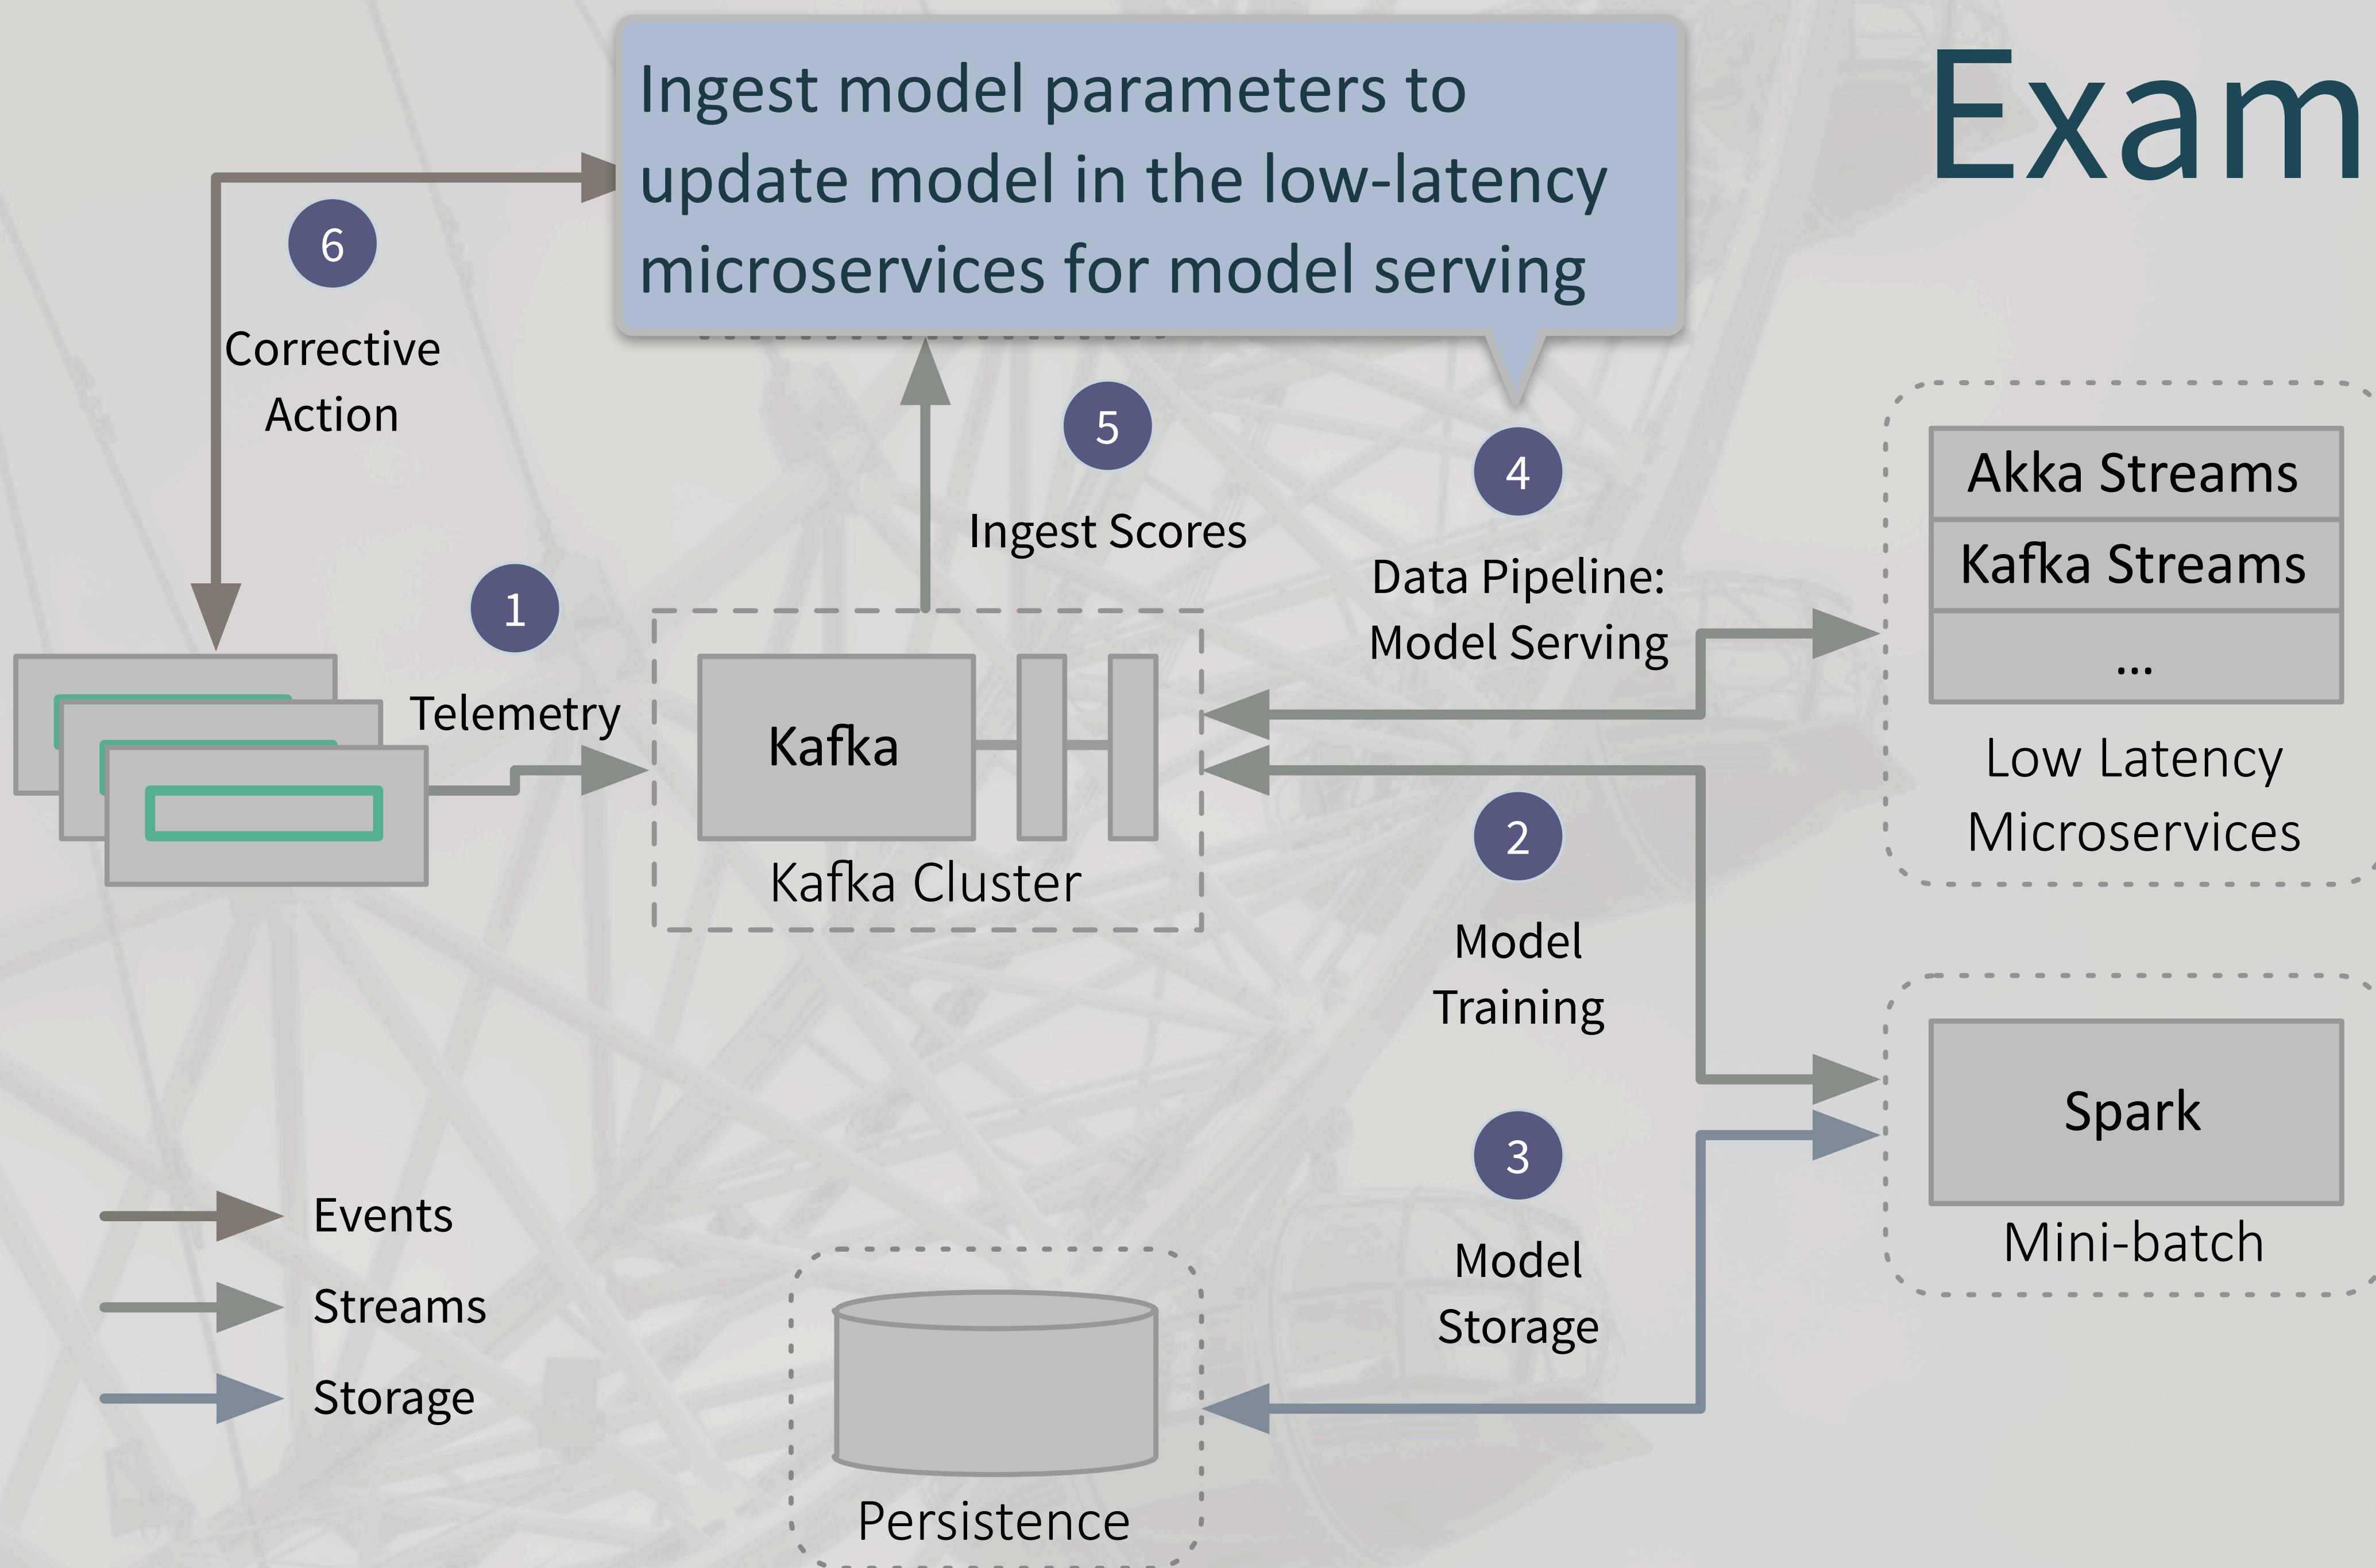
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



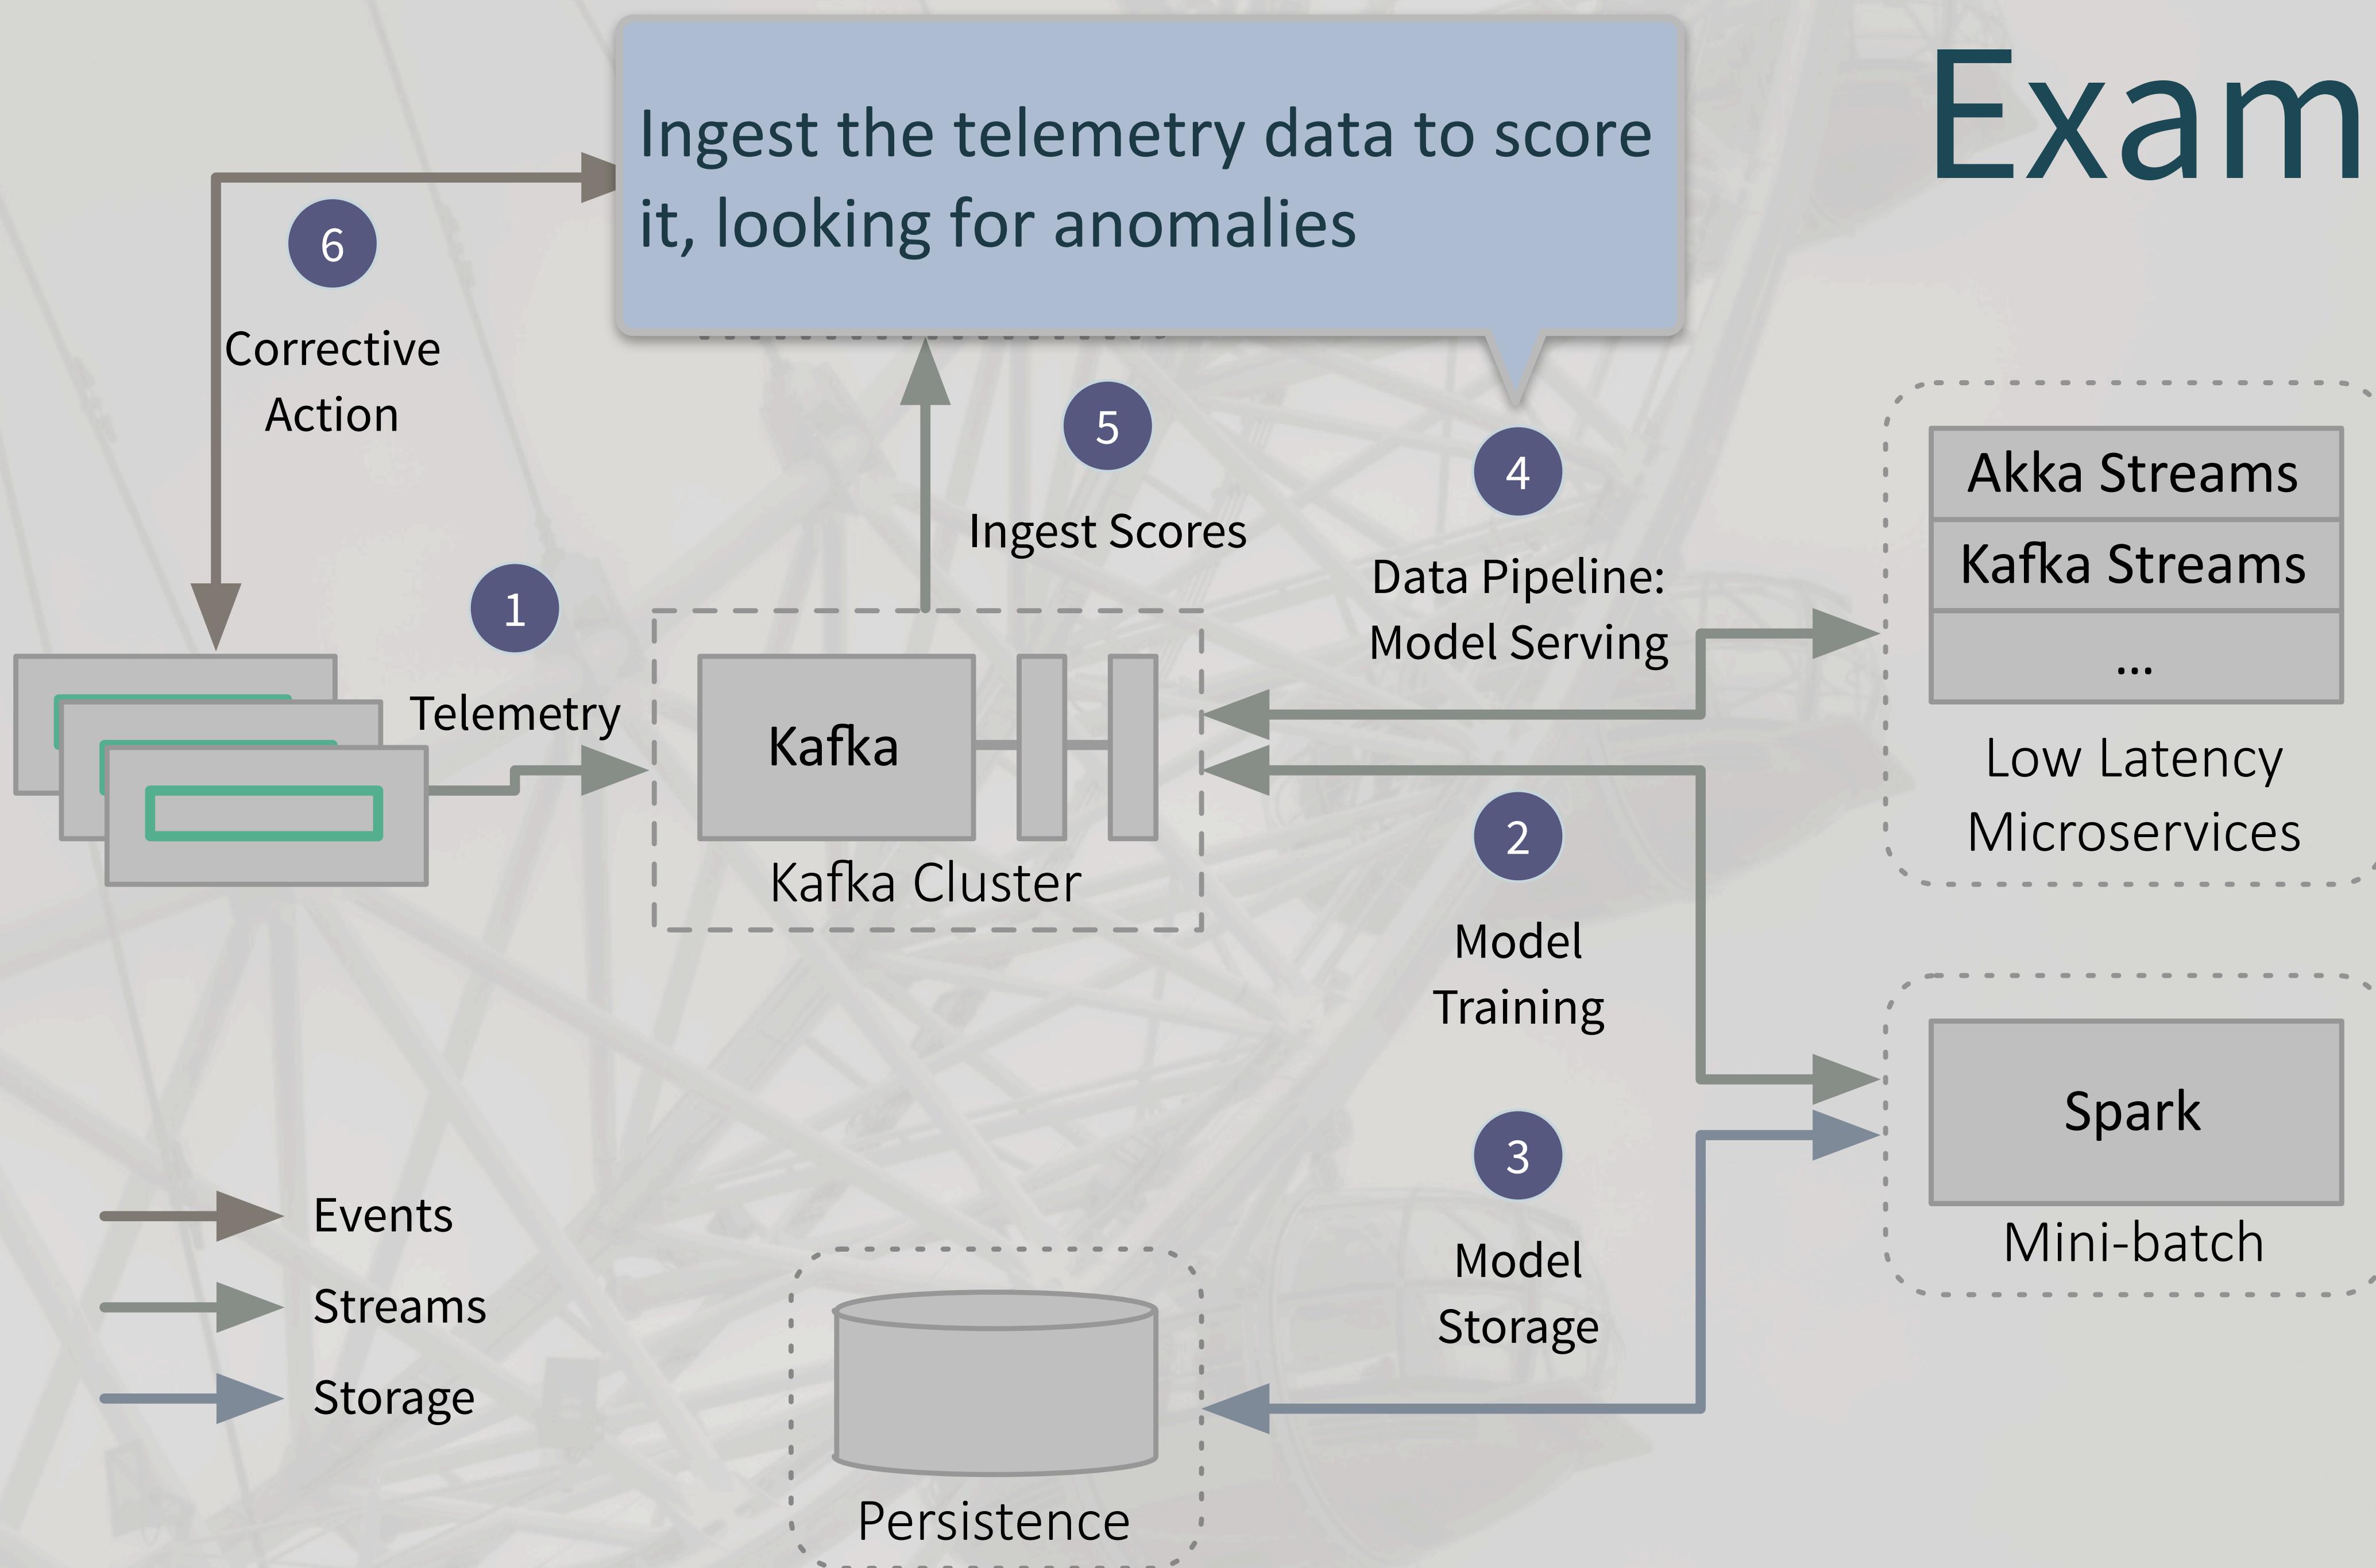
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



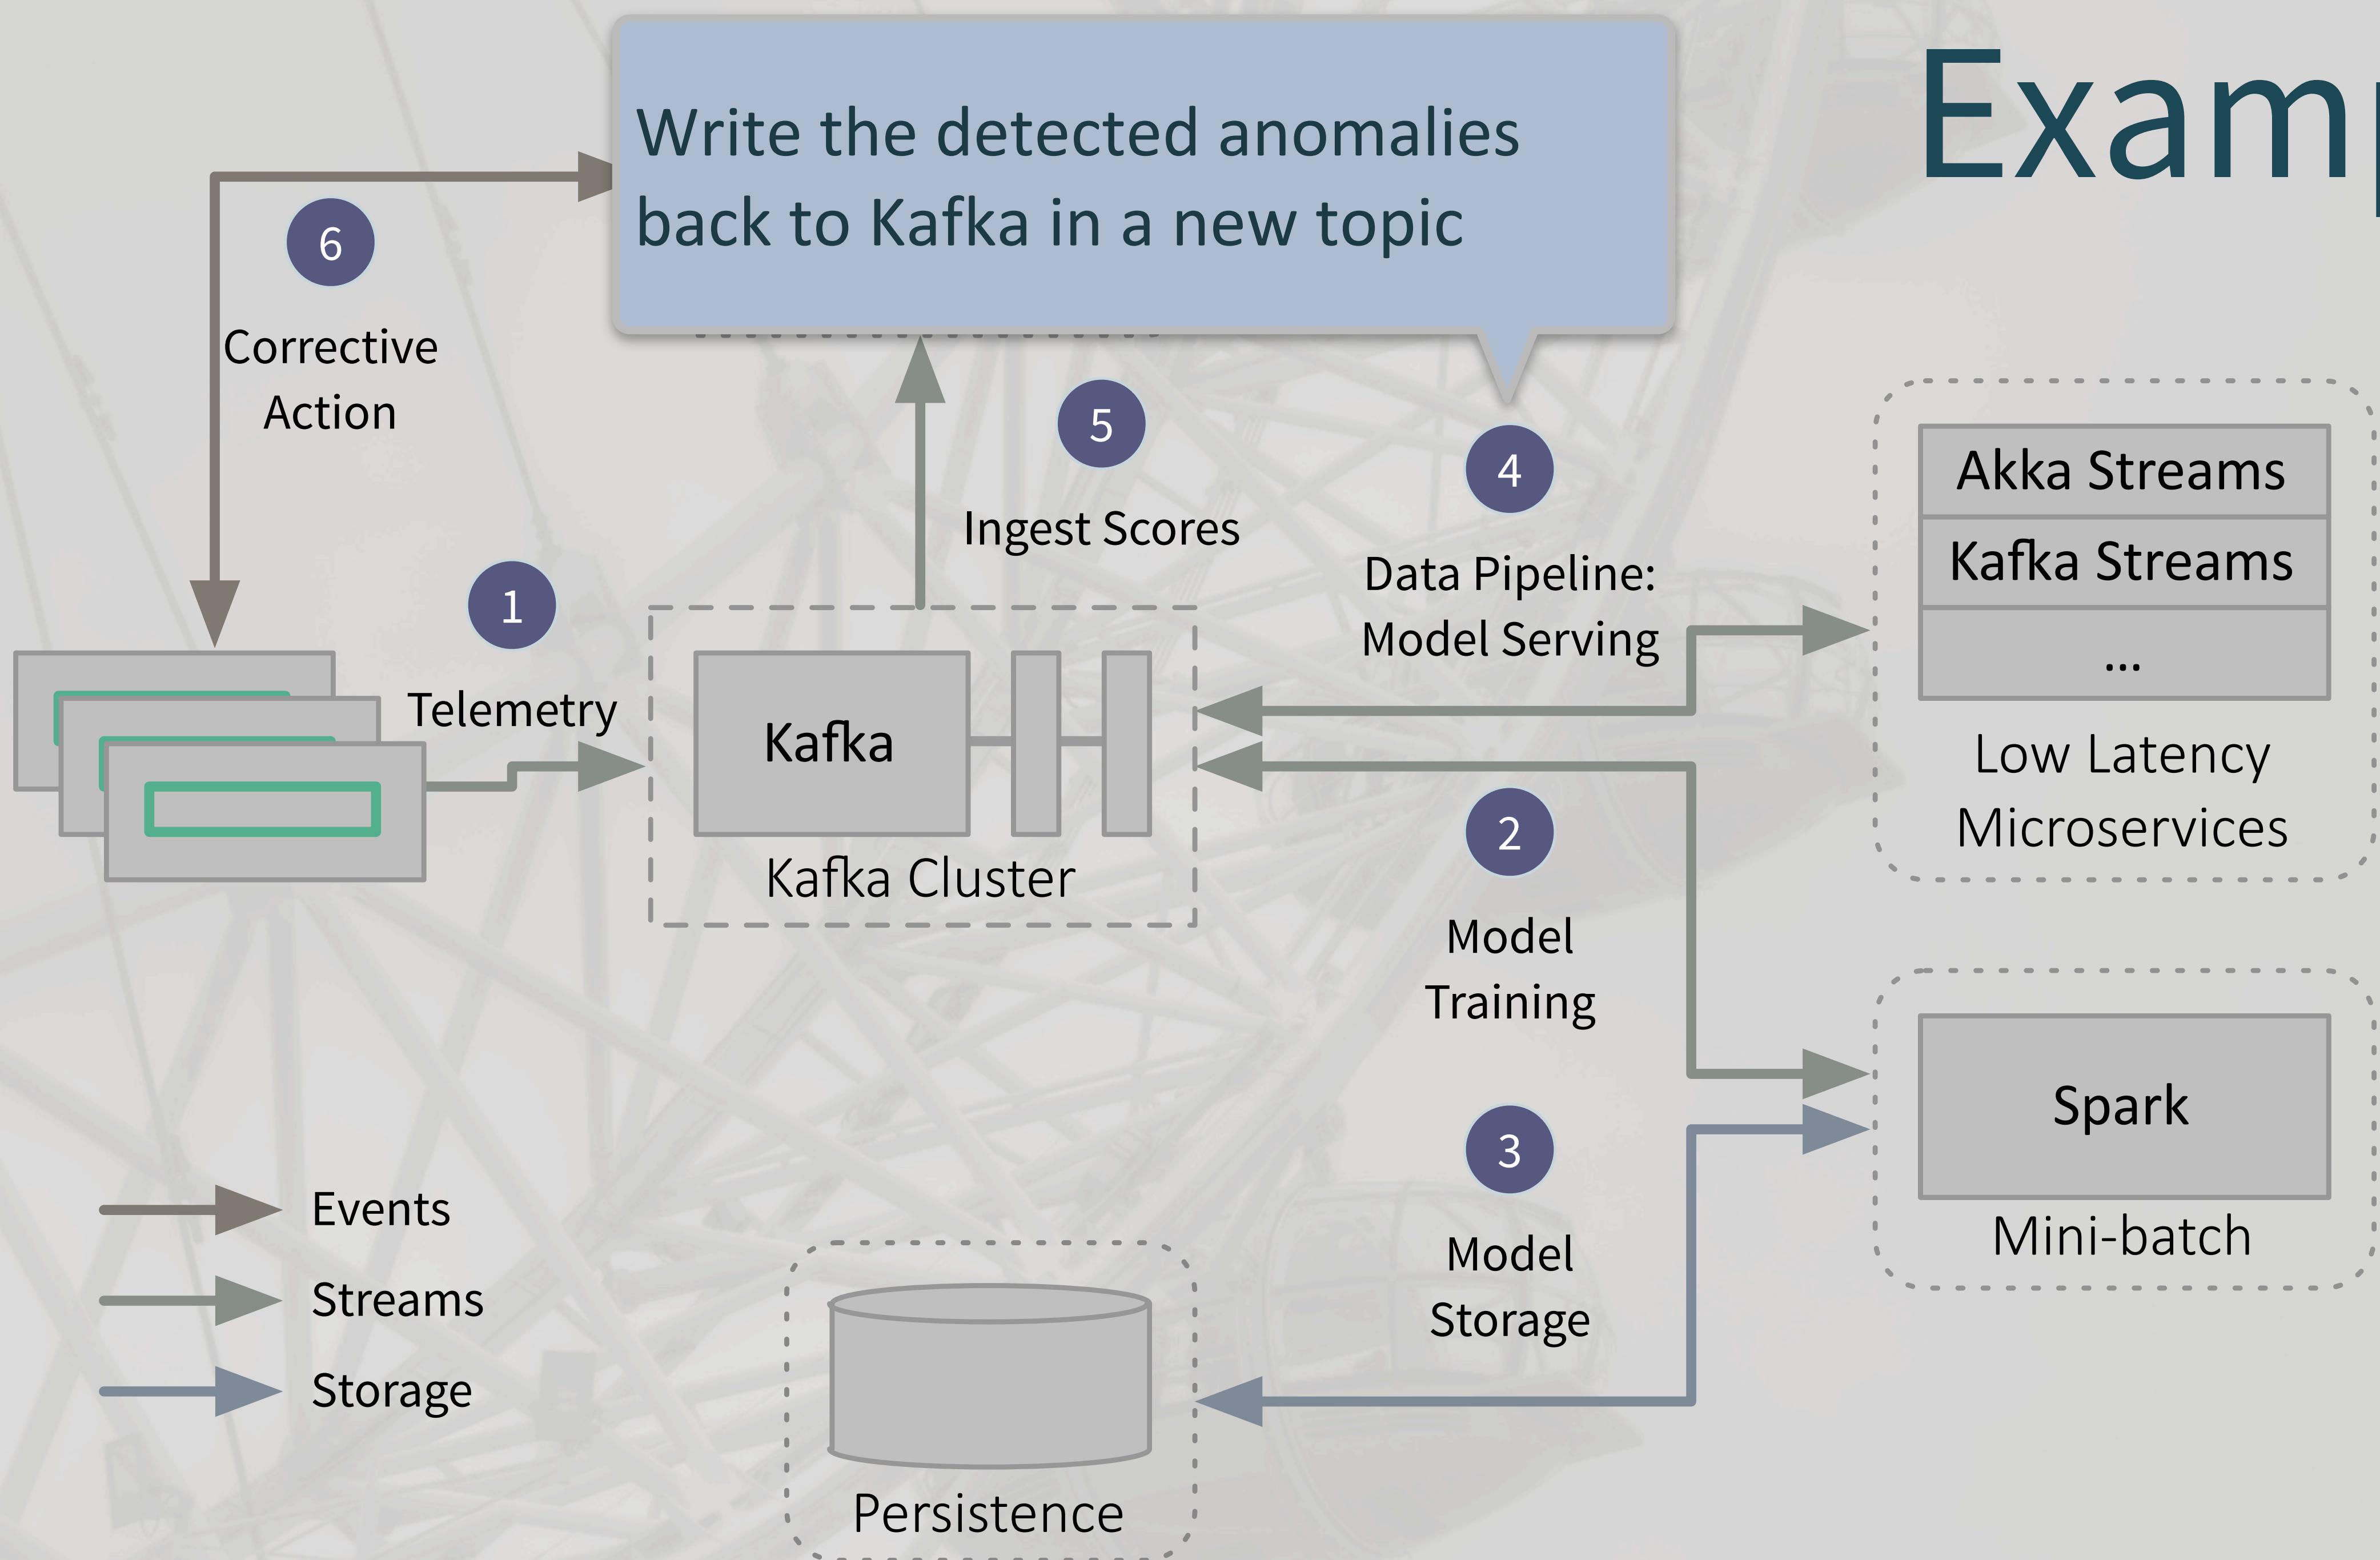
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



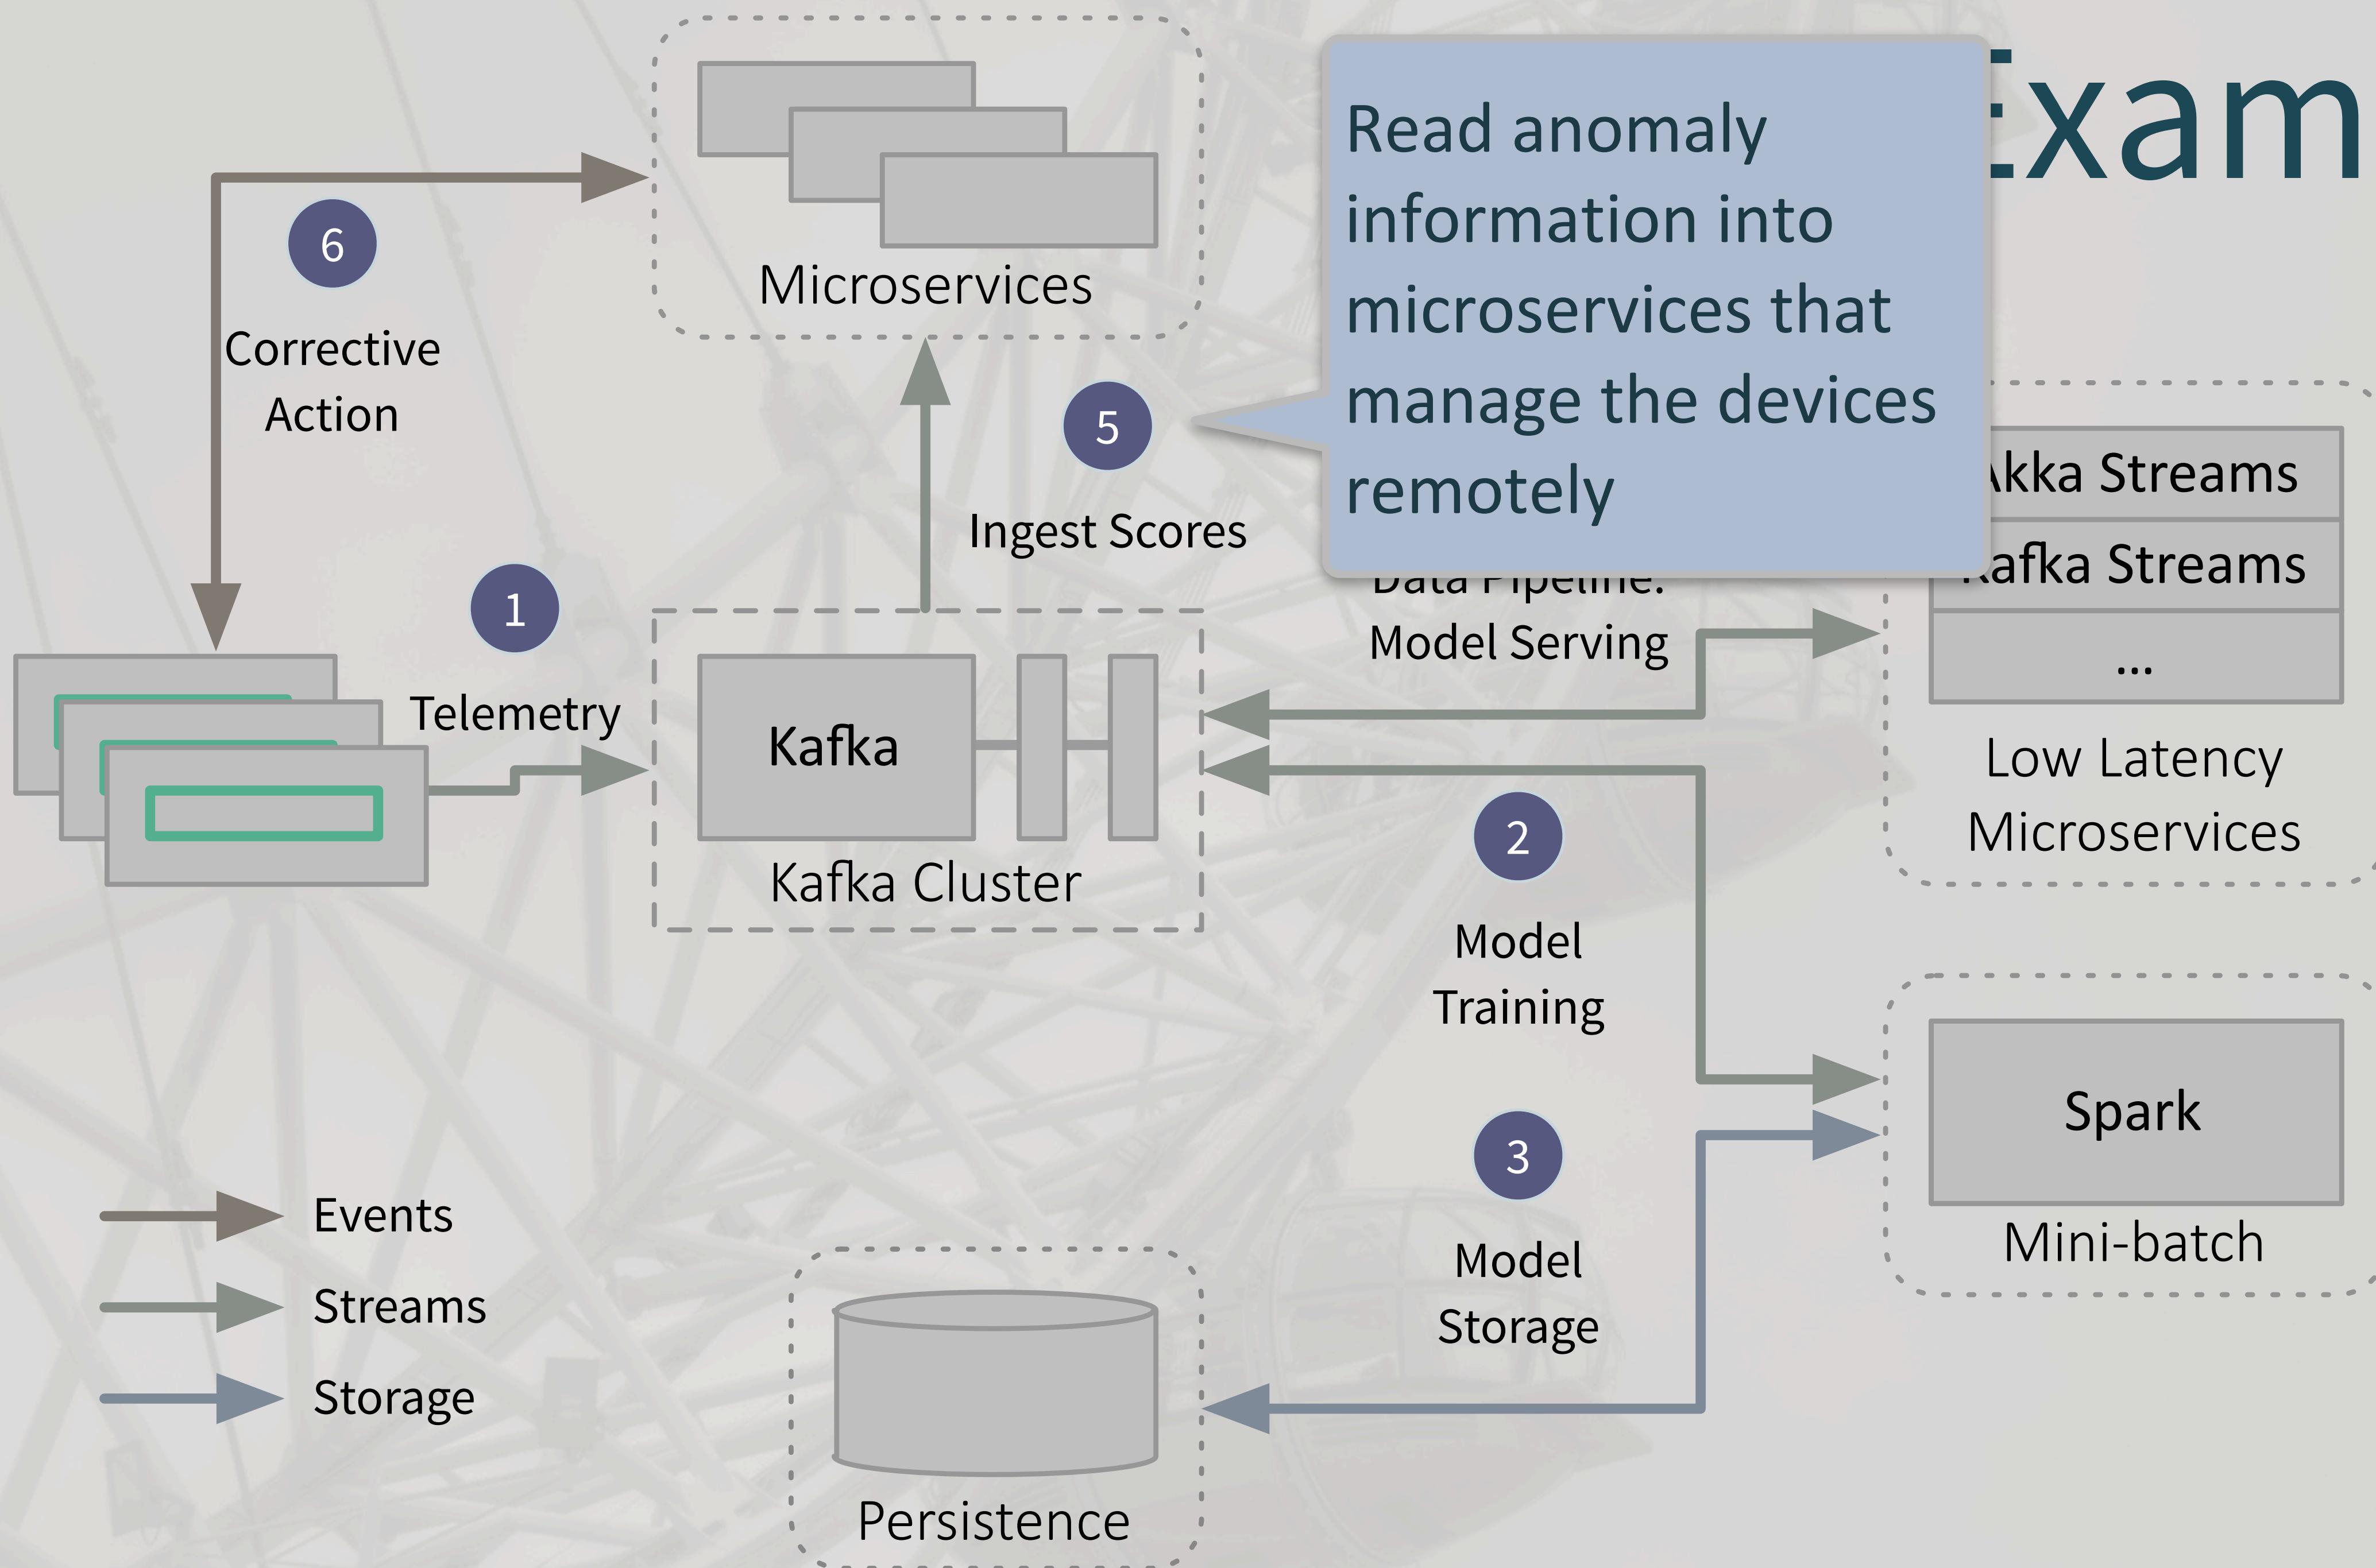
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



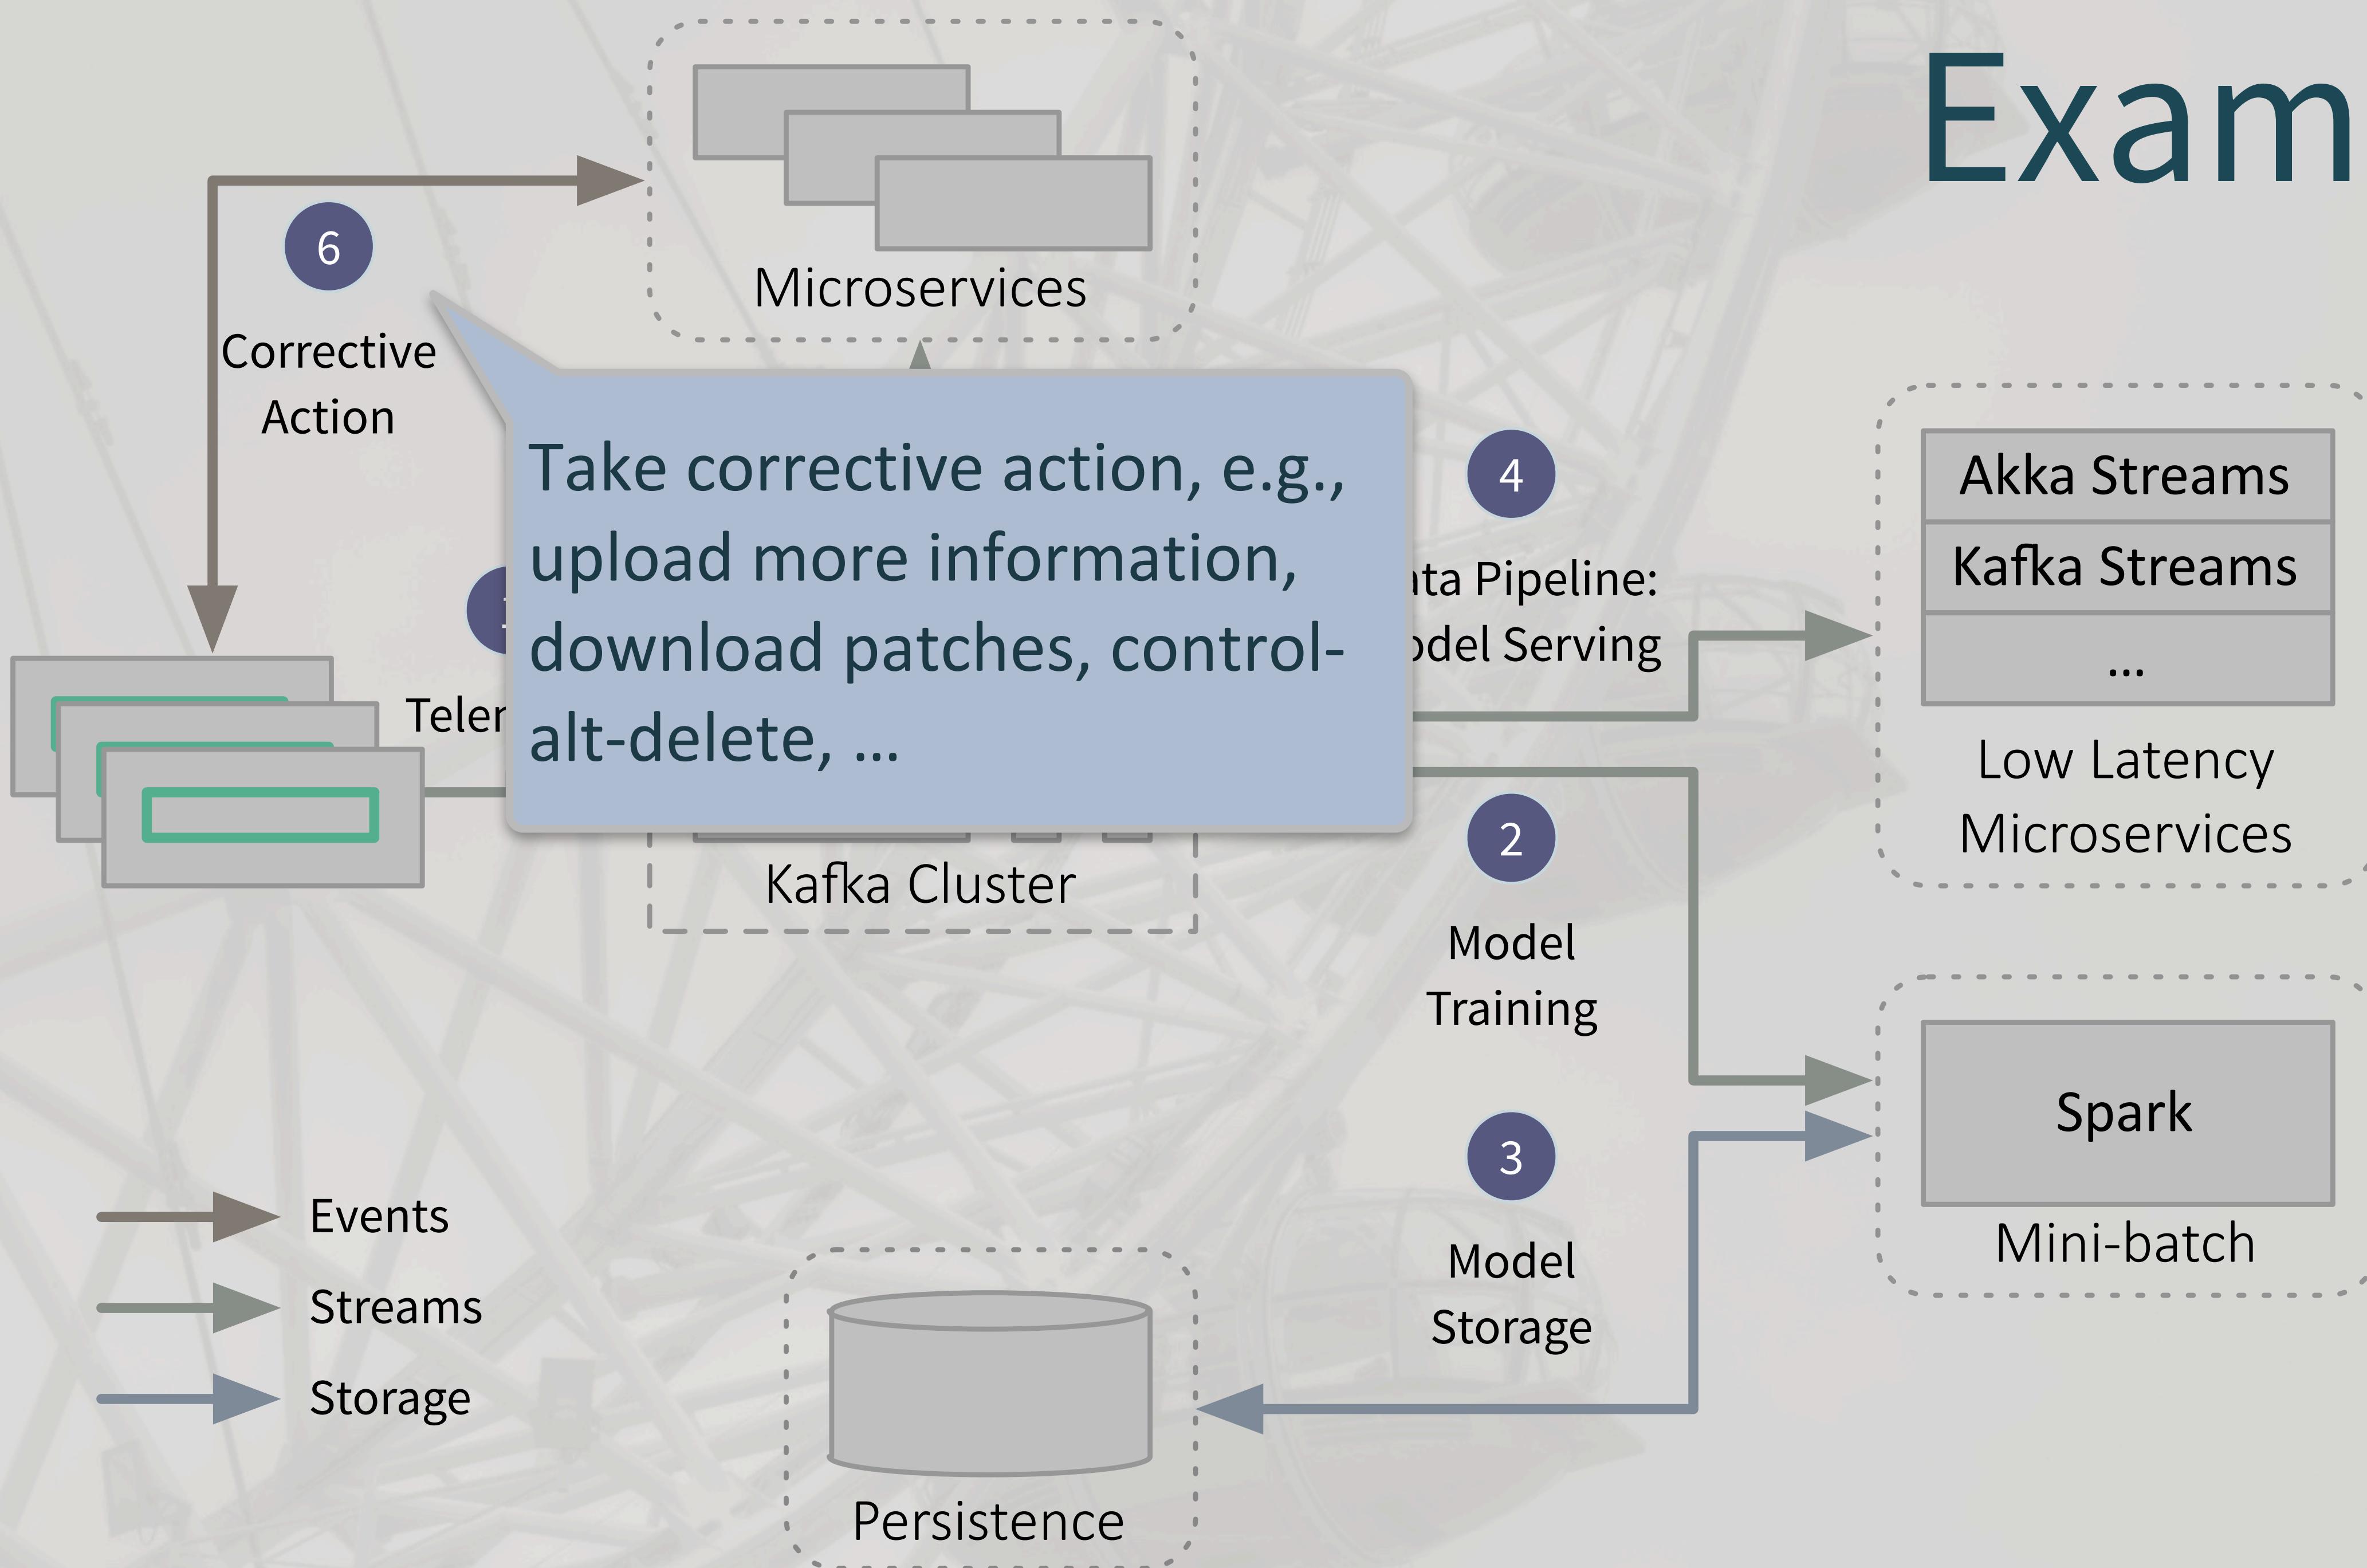
Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Example



Example of how you might implement the HPE use case, where device telemetry is ingested into Kafka, then consumed by a Spark mini-batch or periodic batch job to train models for anomaly detection. The models are written back to a new Kafka topic. Meanwhile, the last-trained model is ingested by low-latency microservices for serving the model and the data is also ingested here for scoring. Scores for anomalies are written to a new topic where other microservices that manage the devices take corrective action.

Real-time Personalization



- Model users to provide real-time marketing based on behavior, location, inventory levels, product promotions, etc.

Real-time Financial



- Drive better business outcomes through real-time risk, fraud detection, compliance, audit, governance, etc.



Internet of Things

- Real-time consumer and industrial device and supply chain management at scale

Legacy Modernization



- Accelerate decision making processes and optimize infrastructure costs by moving from batch to streaming
- Hadoop replacement

And One More... Hard Real Time

SPACE



One more we didn't mention, but for completeness... there is "hard" real-time, like flight control software. This is beyond our scope here, because it usually involves custom hardware, real-time operating systems, and software written in C/C++, all of which squeeze out as much latency and overhead as possible. However, if you do high-frequency trading, then this is relevant to you ;)

And One More... Hard Real Time

SPACE

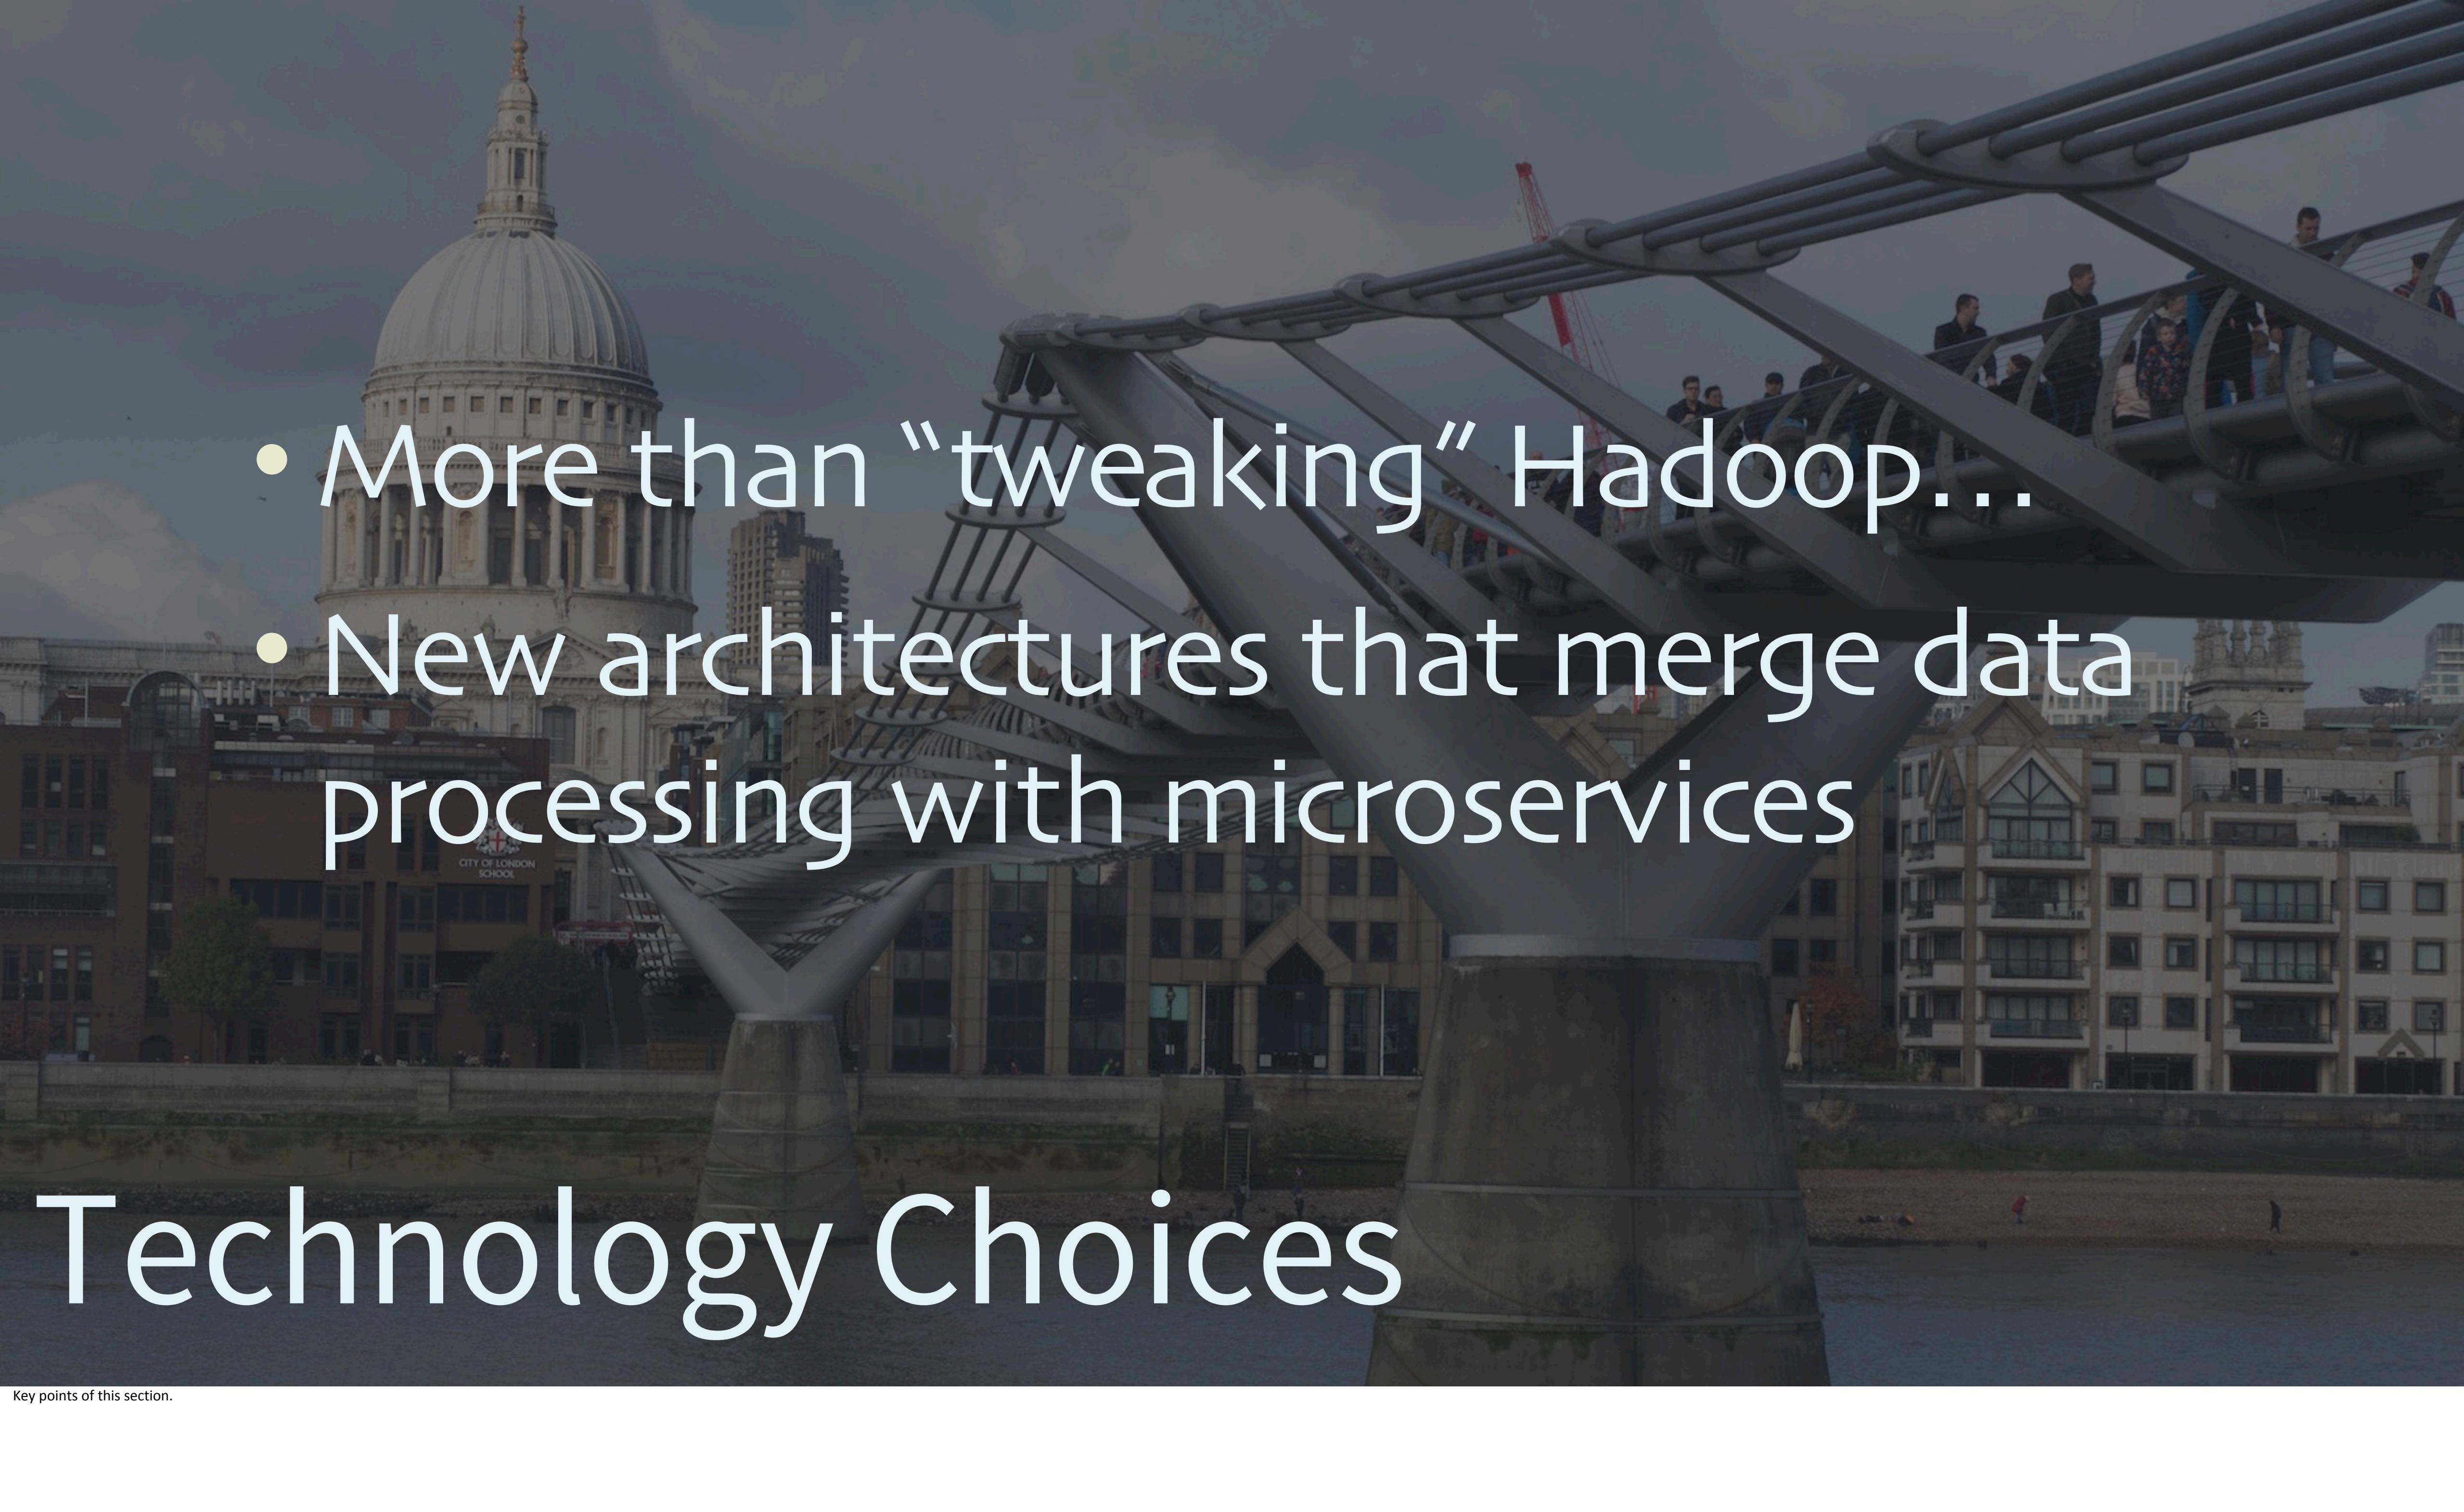
- Flight control software (and high-frequency training) require custom hardware, real-time operating systems, C/C++ software, etc.



Technology Choices

Now let's explore technology choices (ignoring the hard real-time case...)

Photo: Millenium Bridge with St. Pauls in the background

- 
- More than “tweaking” Hadoop...
 - New architectures that merge data processing with microservices

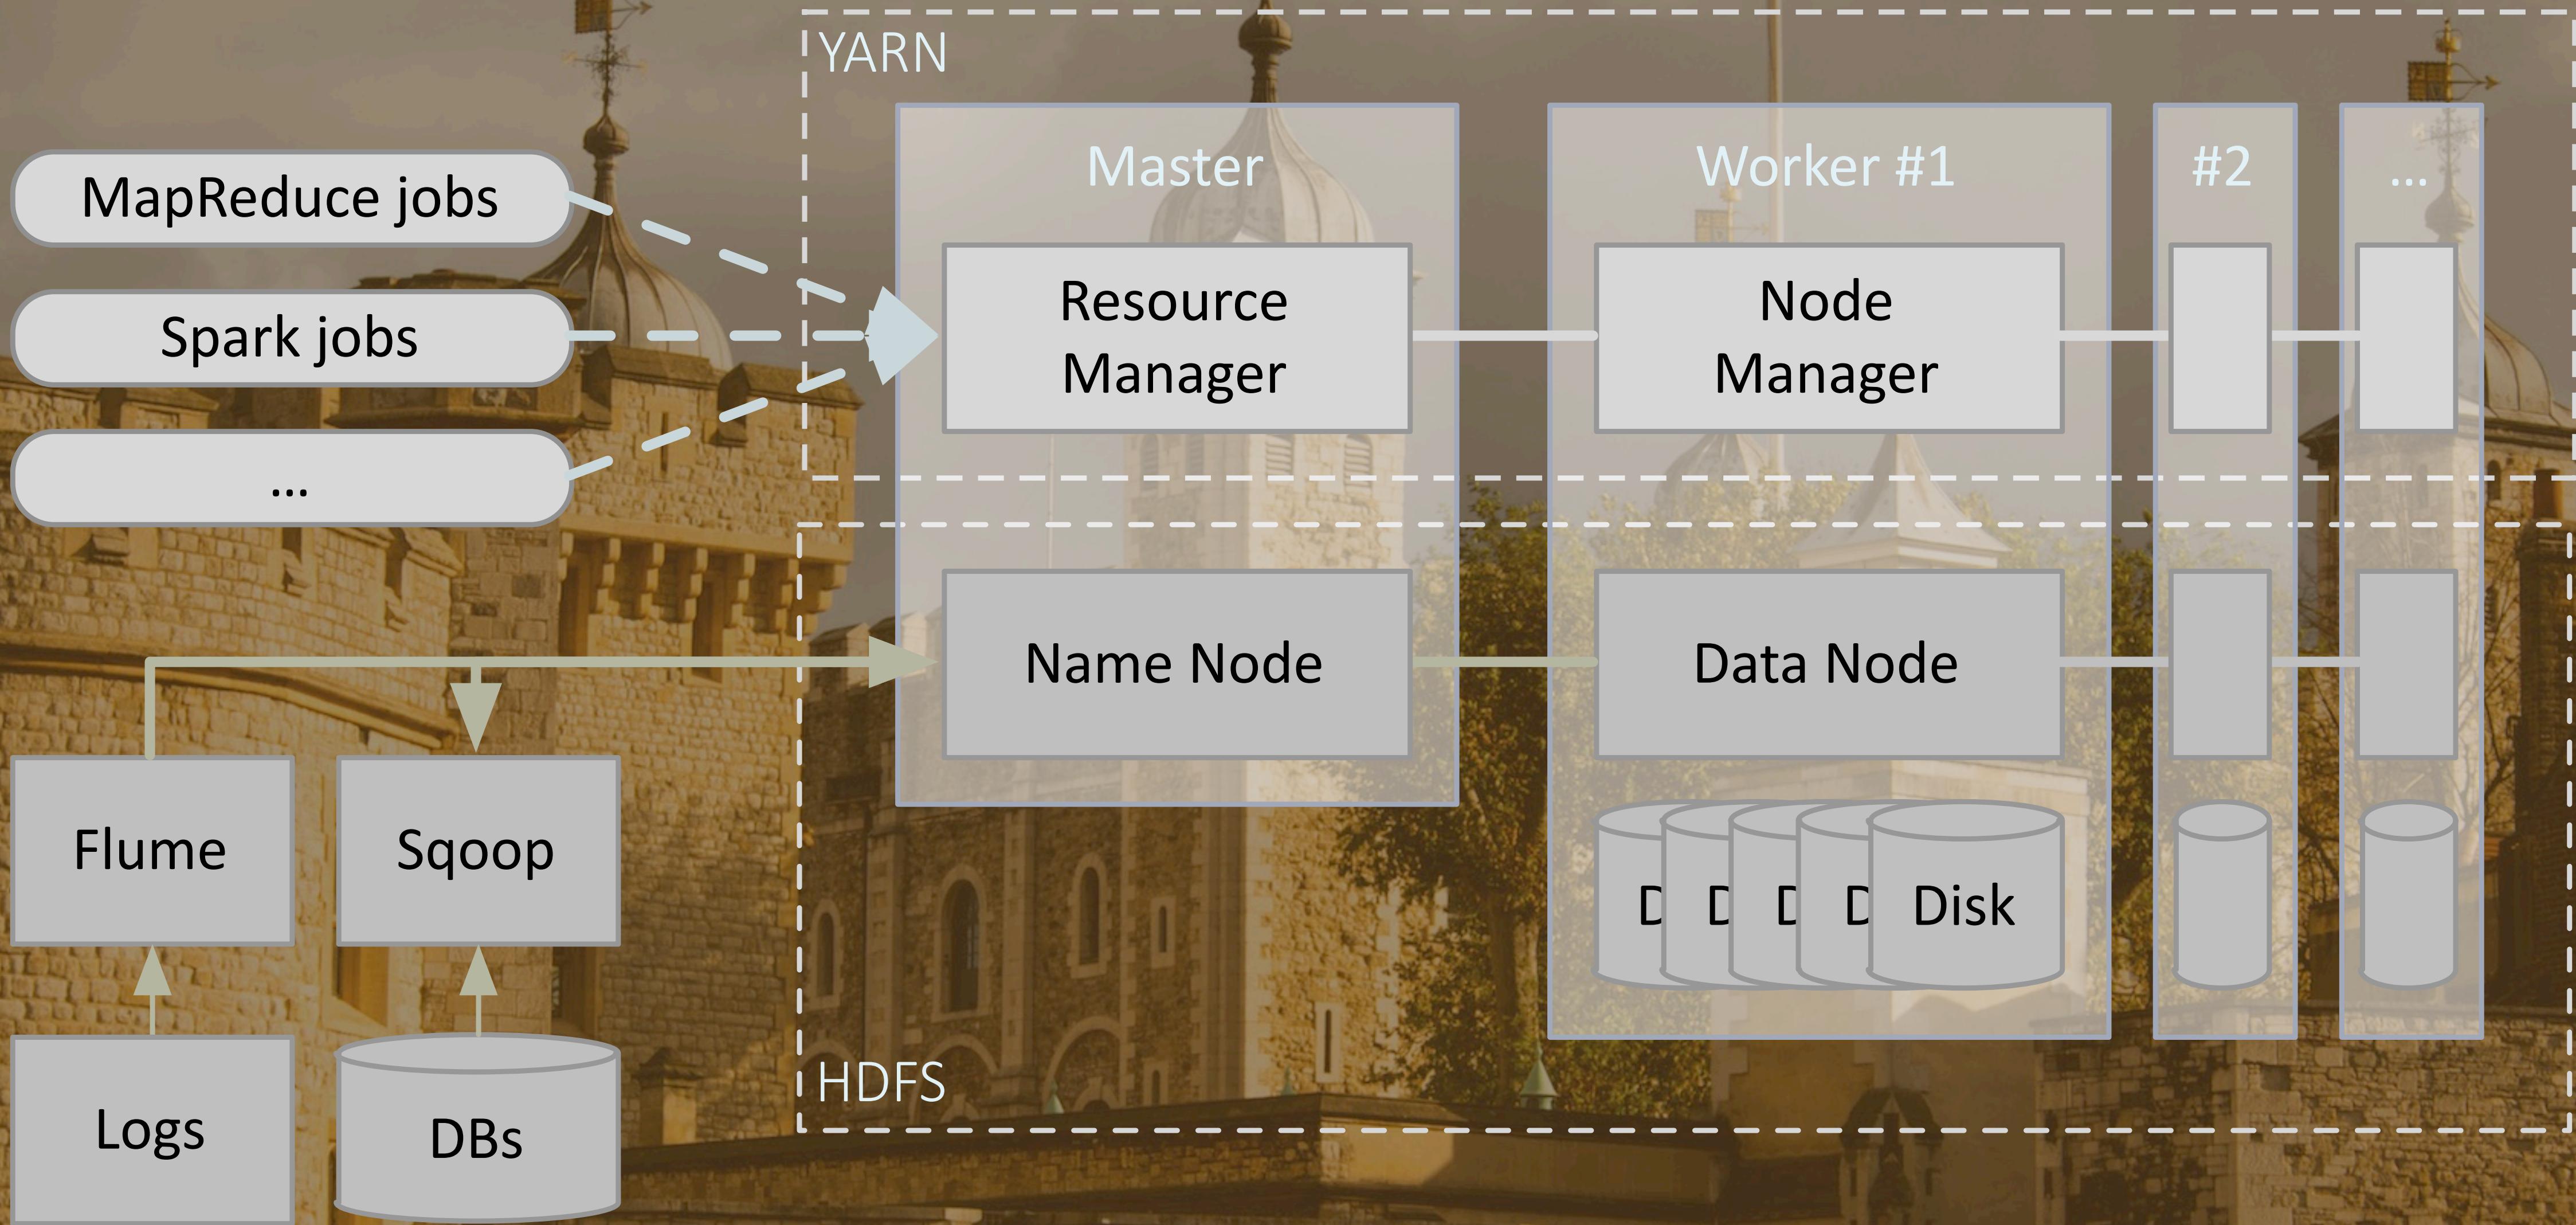
Technology Choices

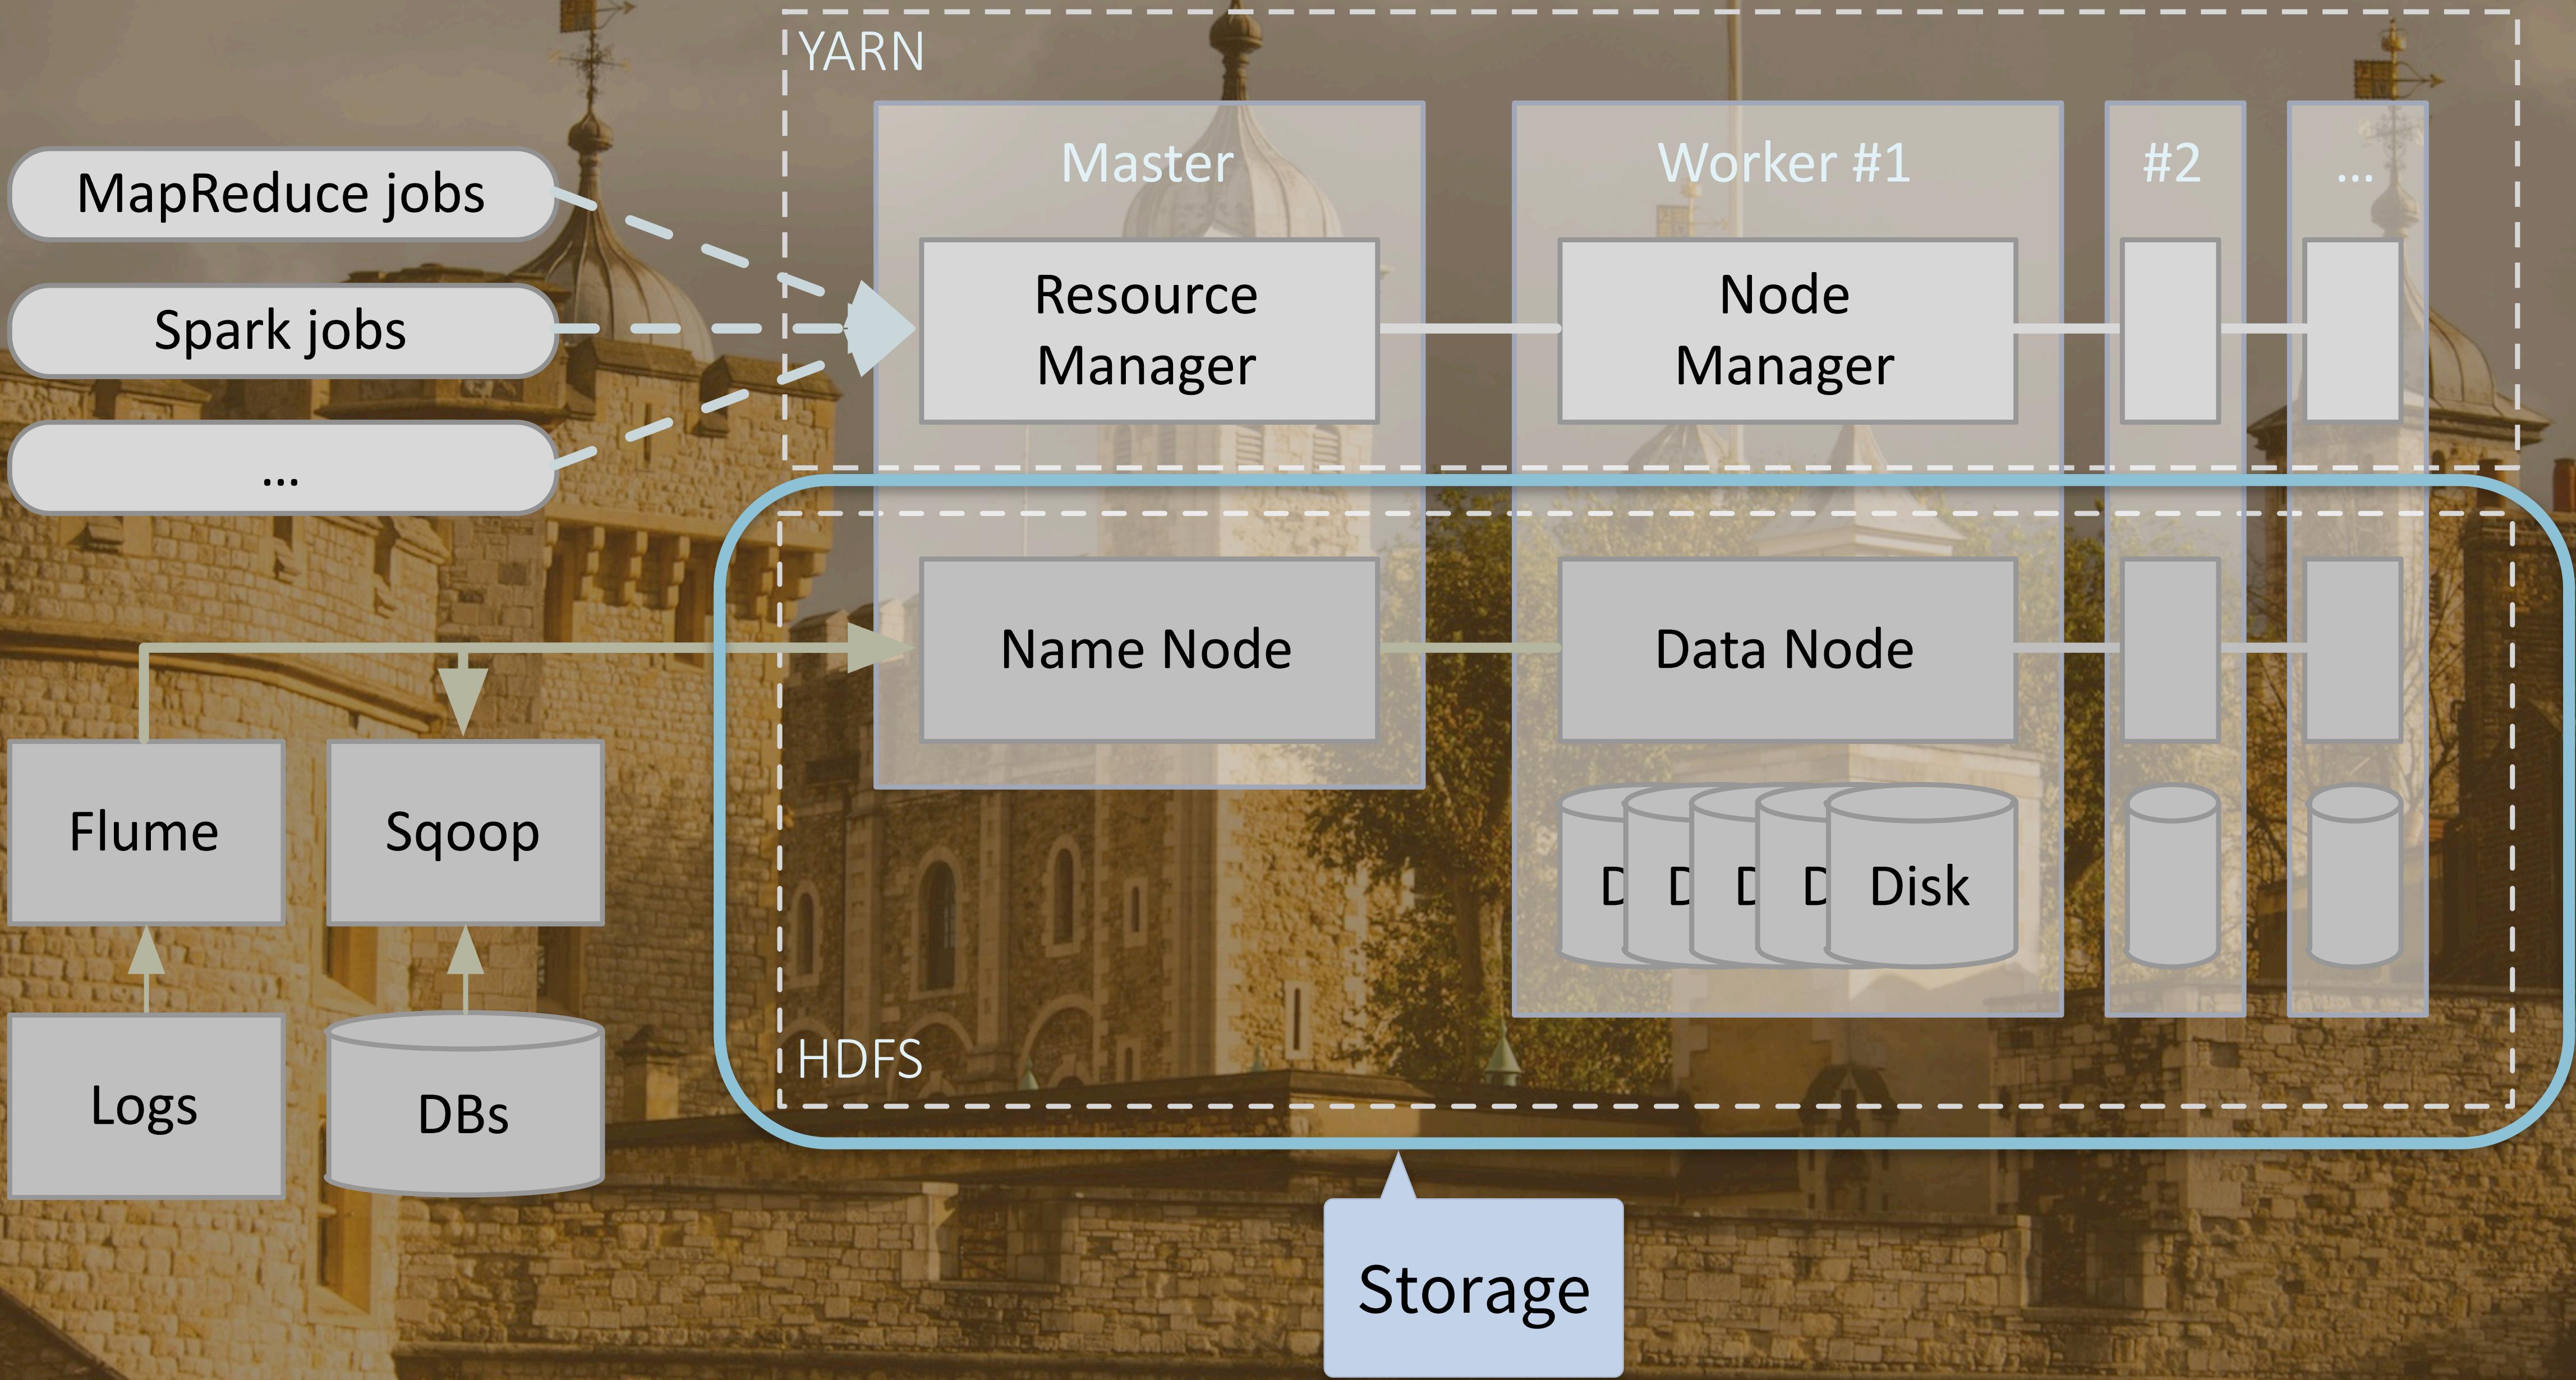
Hadoop



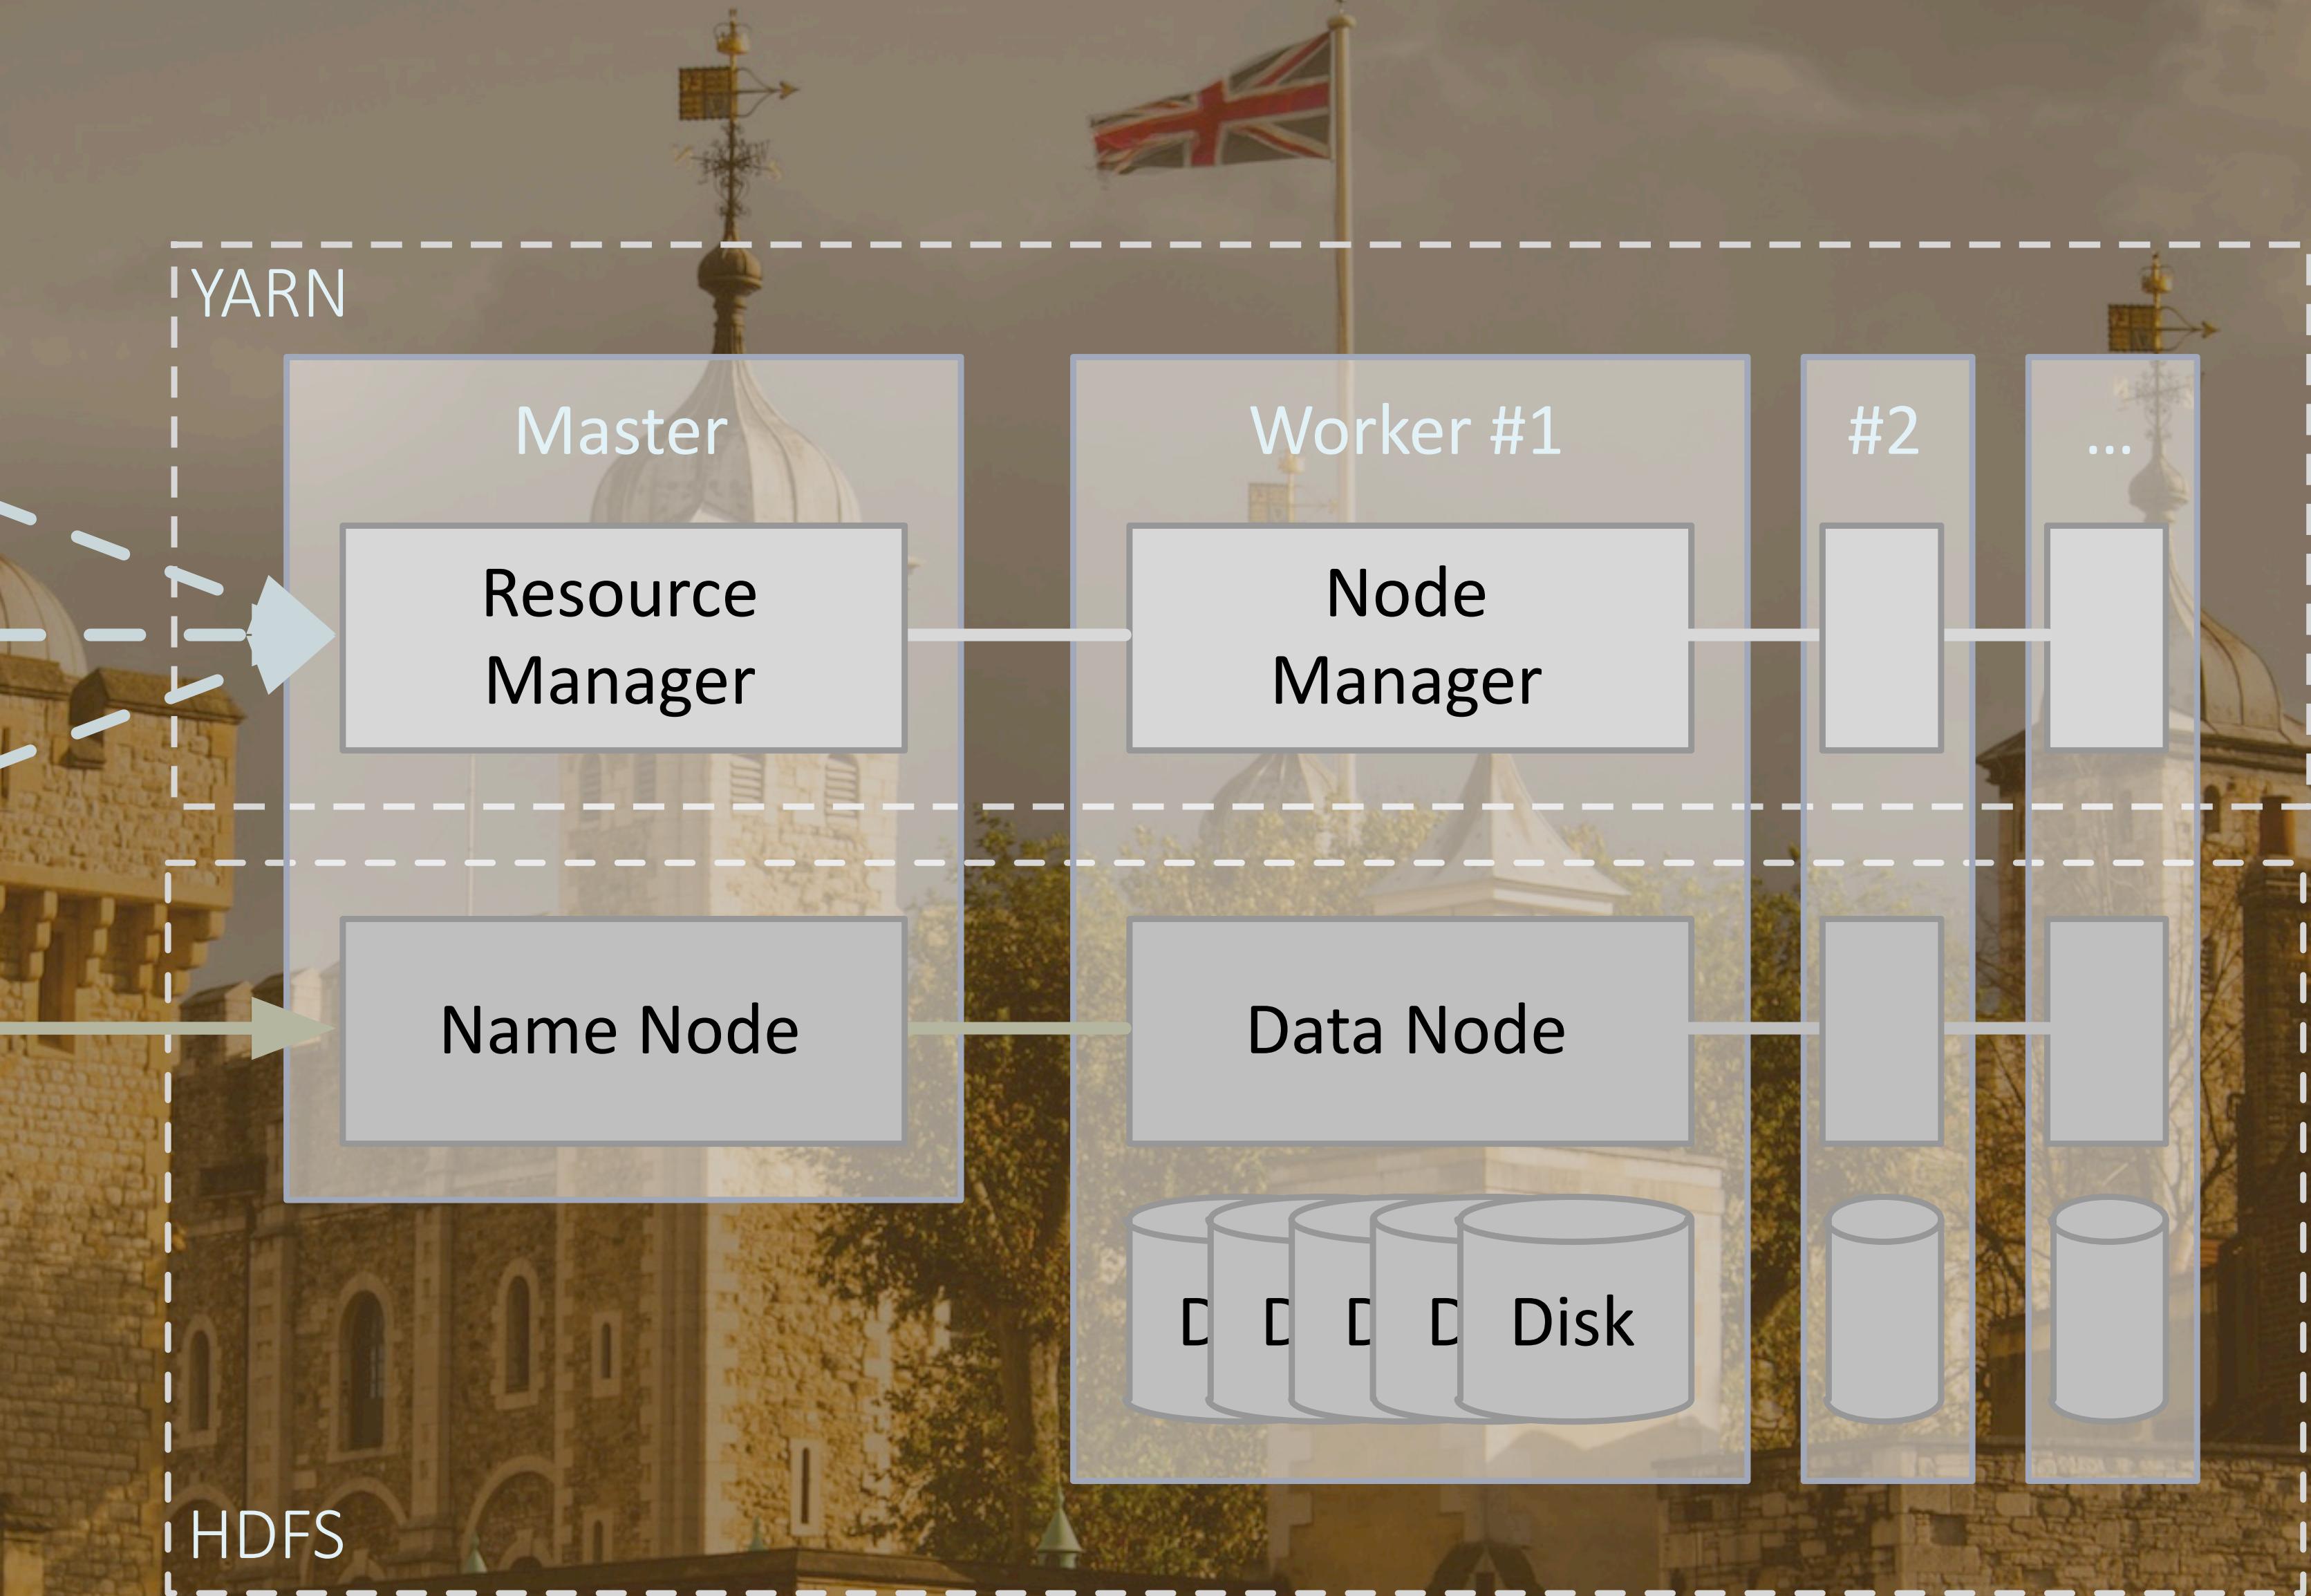
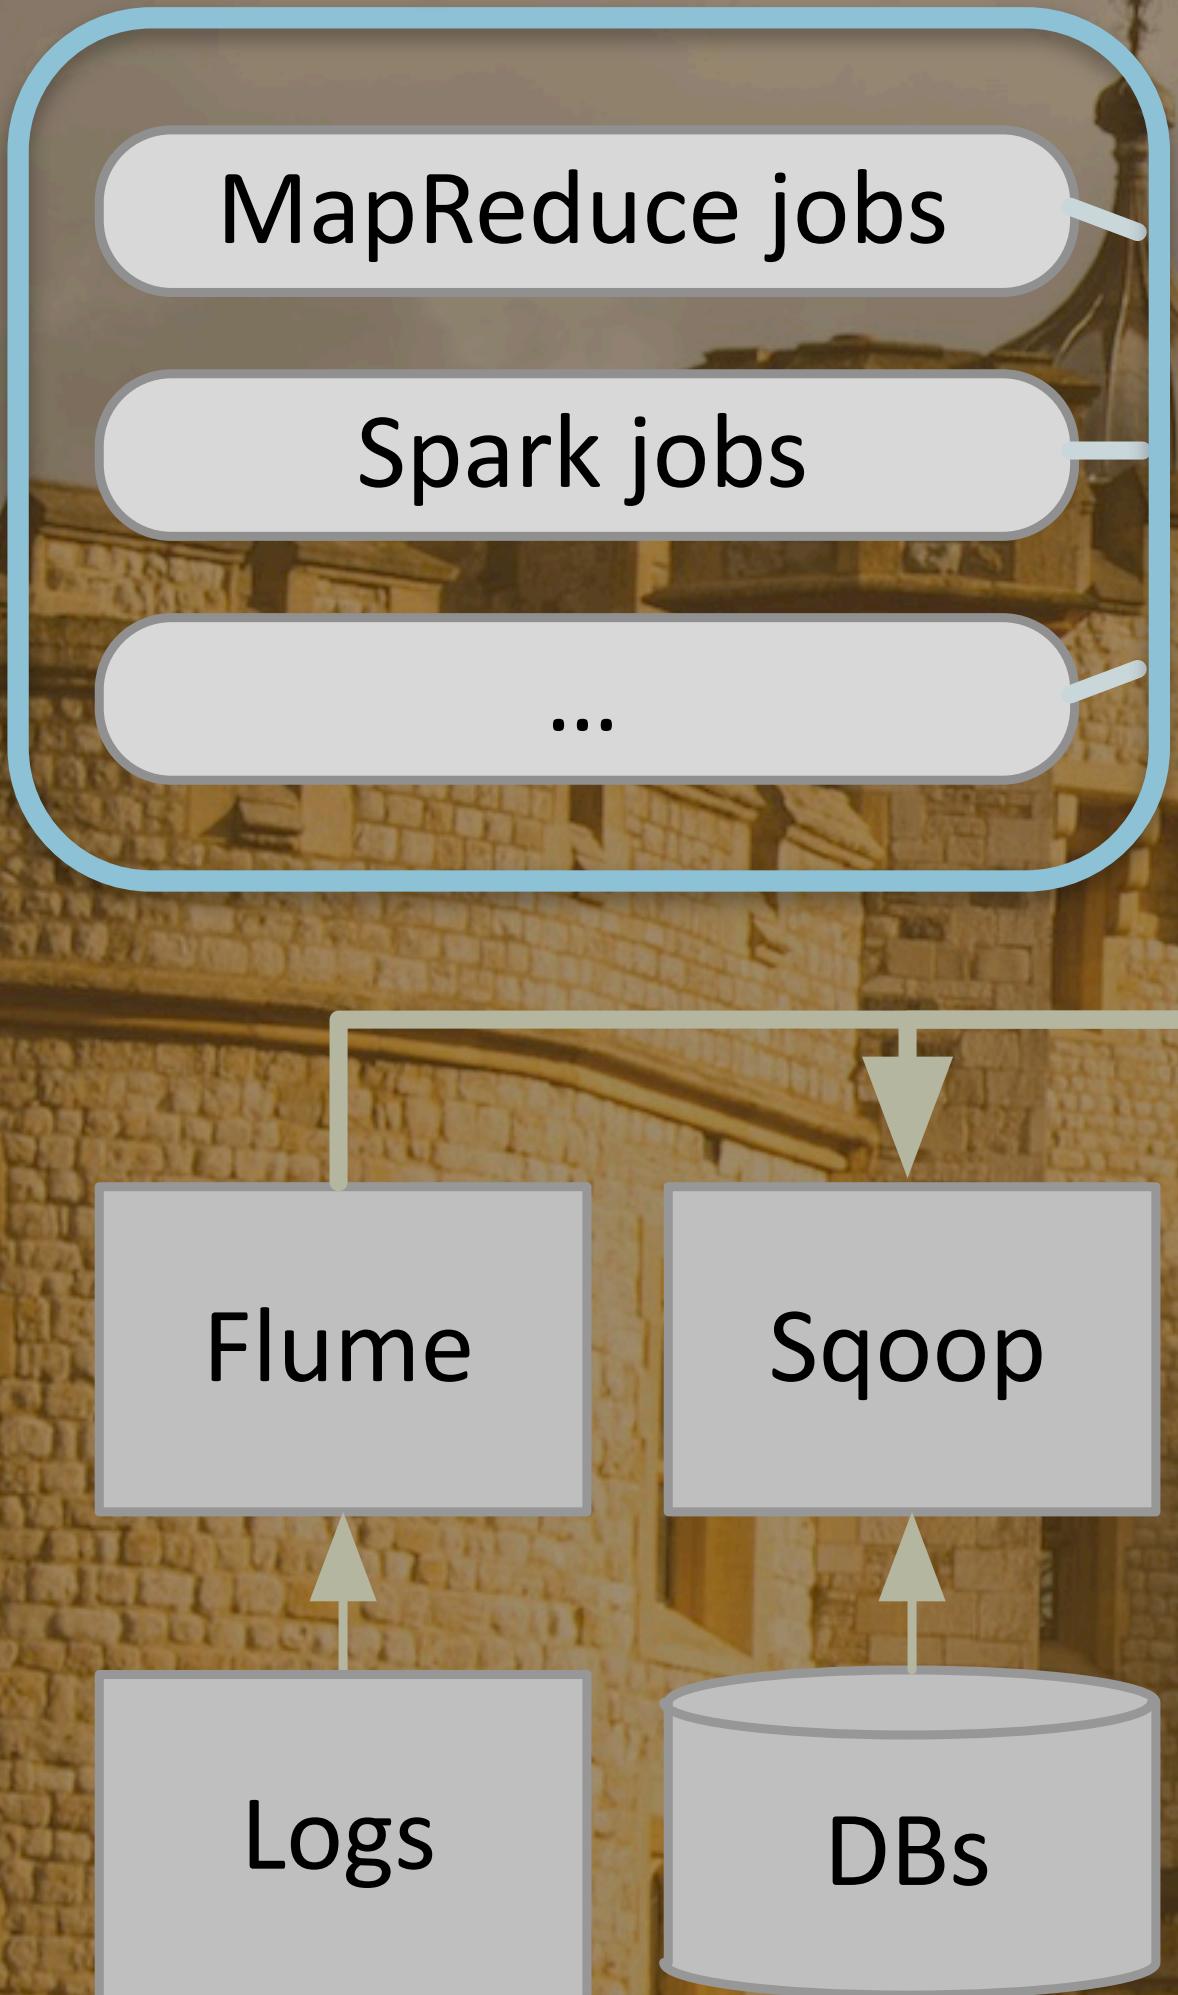
Photo: Detail of the Tower of London

- 
- The background of the slide is a photograph of a historic stone building, likely a university or institutional complex, featuring several towers with domes and weathervanes. A Union Jack flag flies from a pole in front of the building. The sky is overcast.
- Data warehouse replacement
 - Historical analysis
 - Interactive exploration
 - Offline training of machine learning models
 -

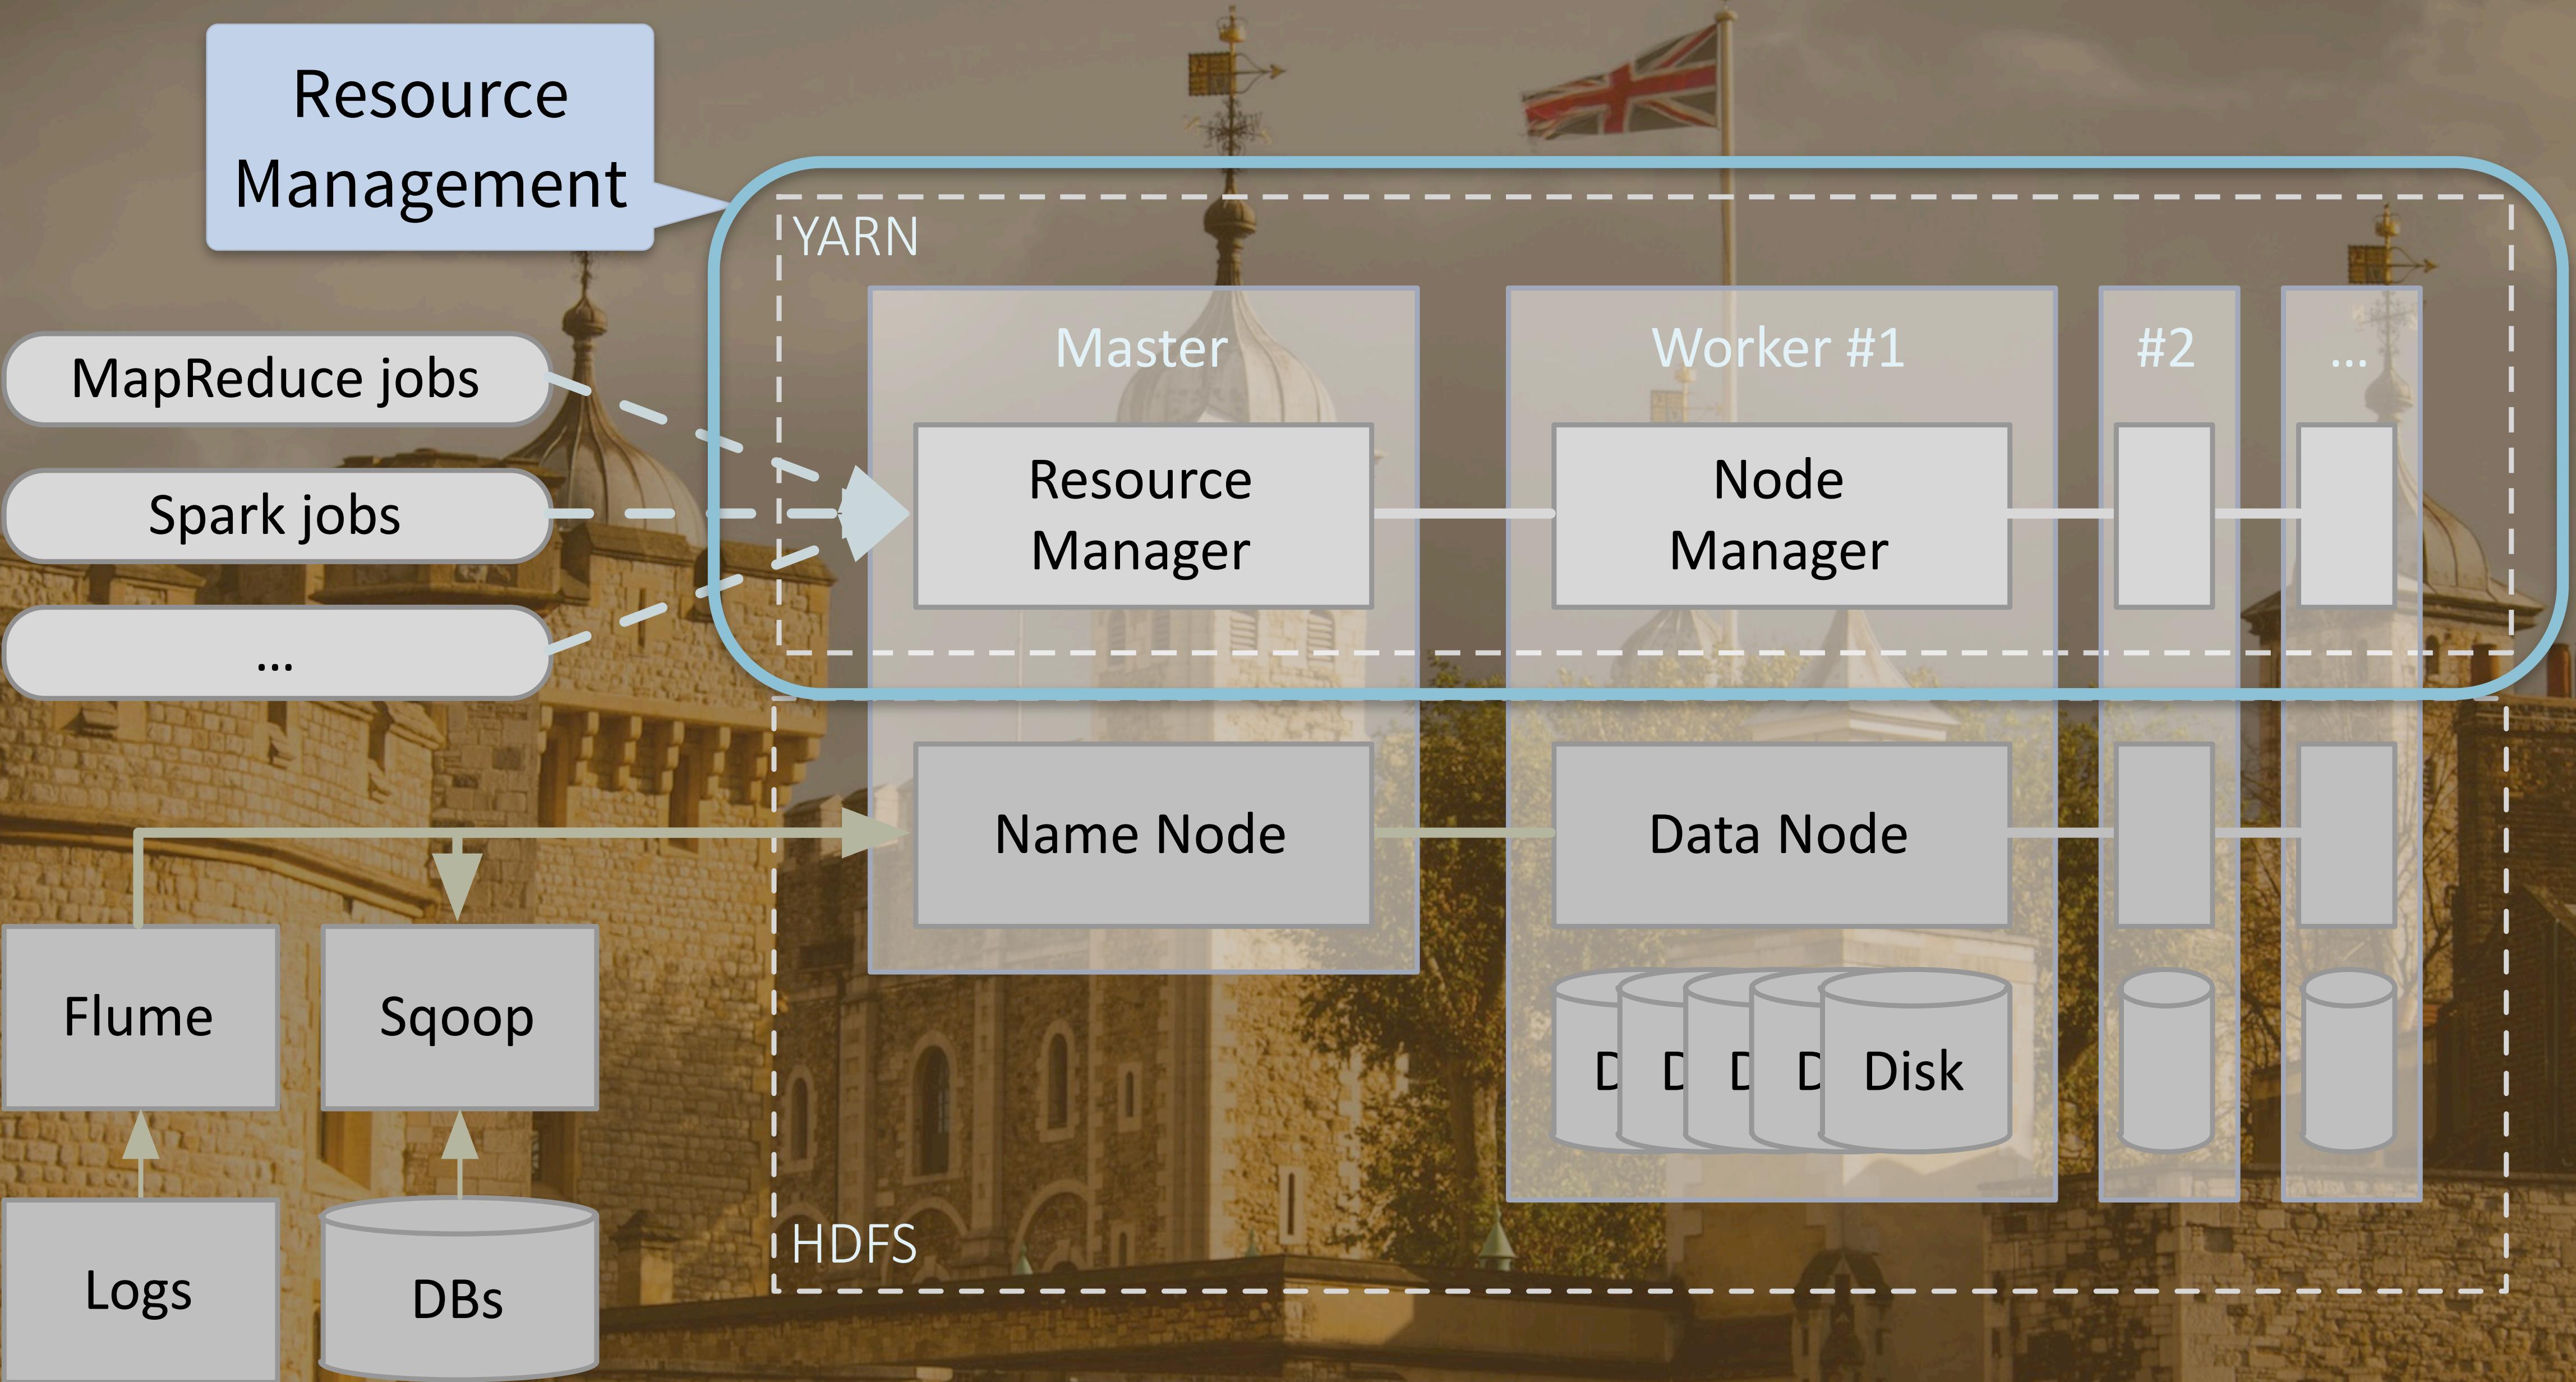




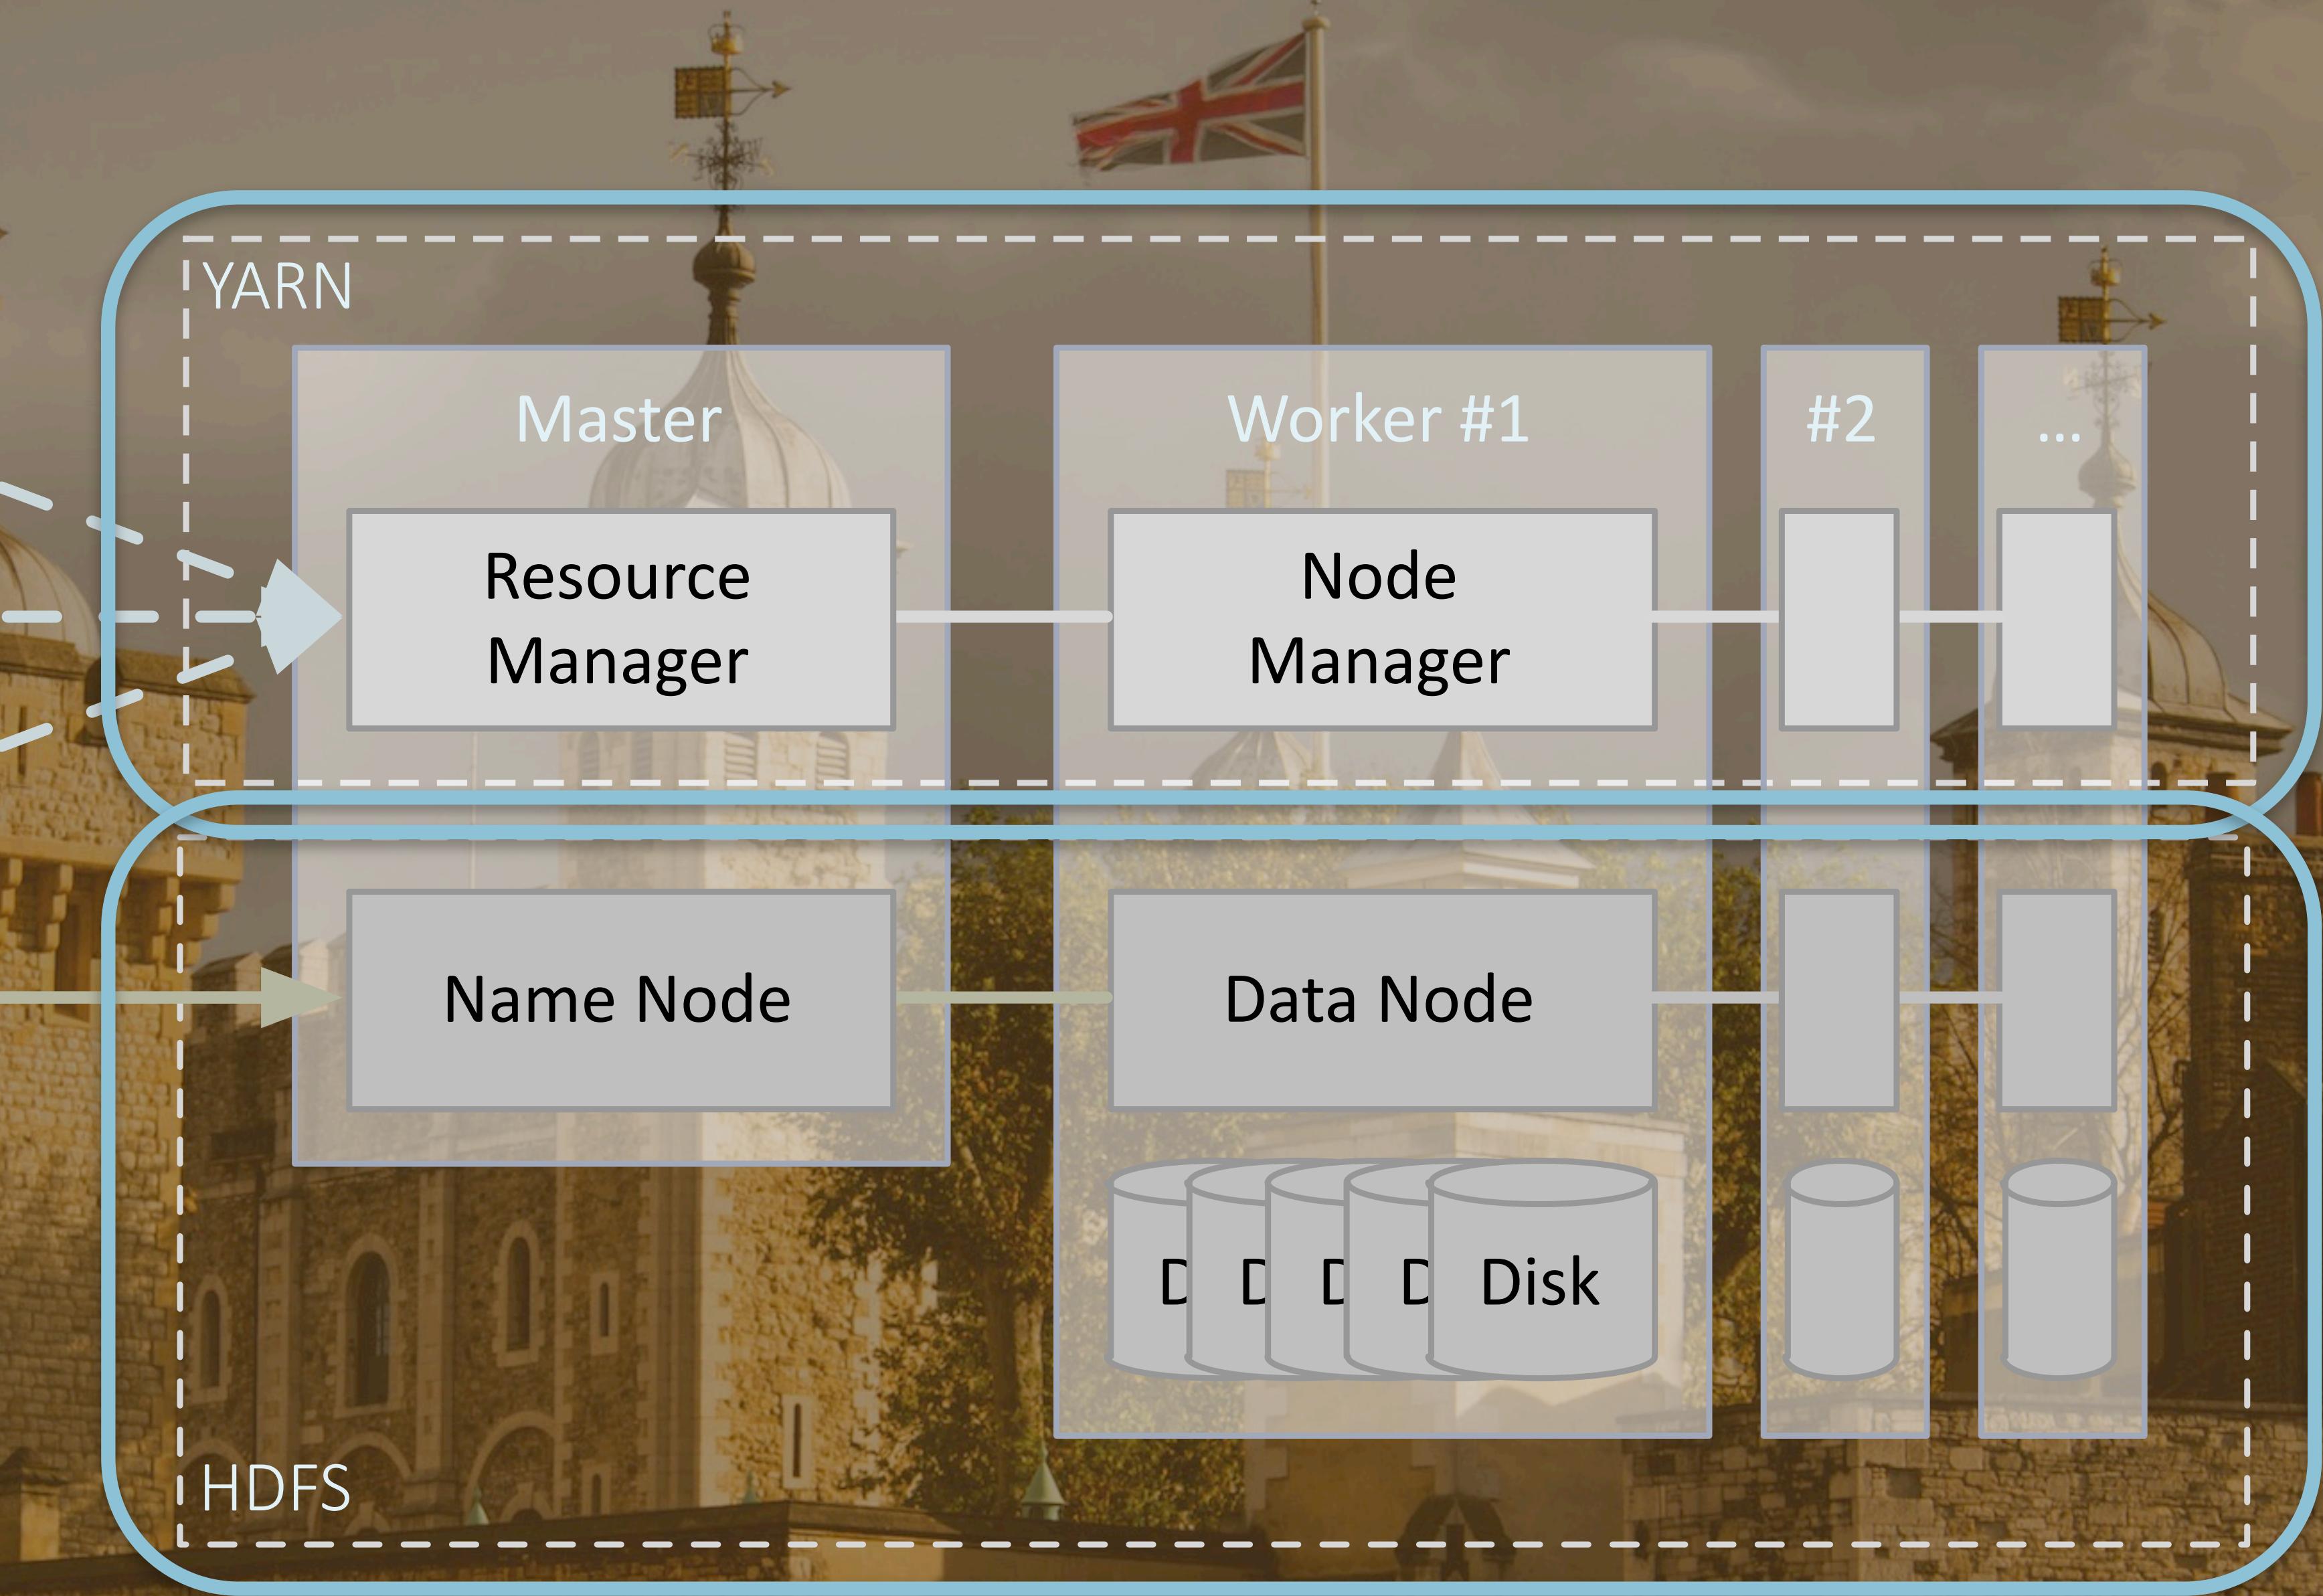
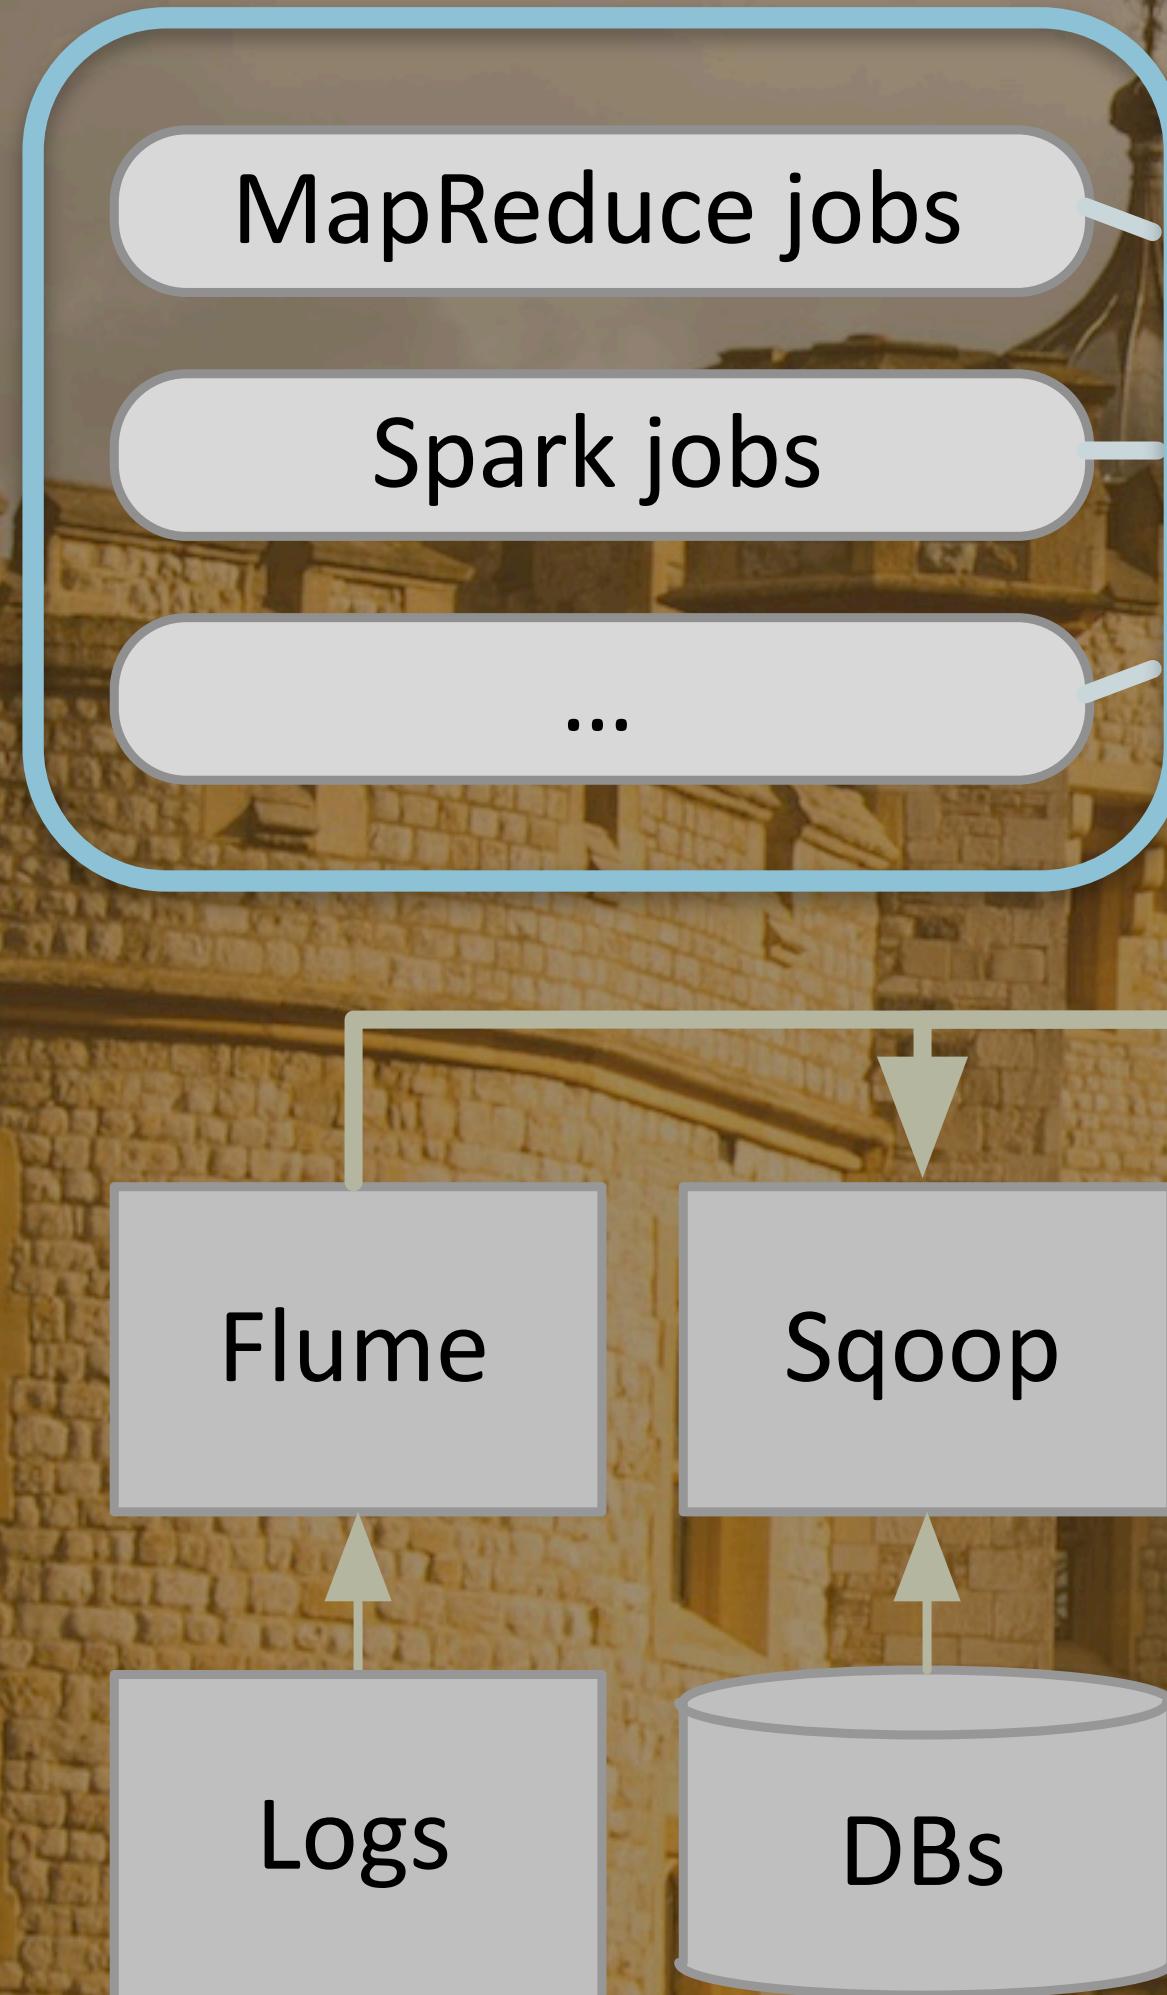
Compute

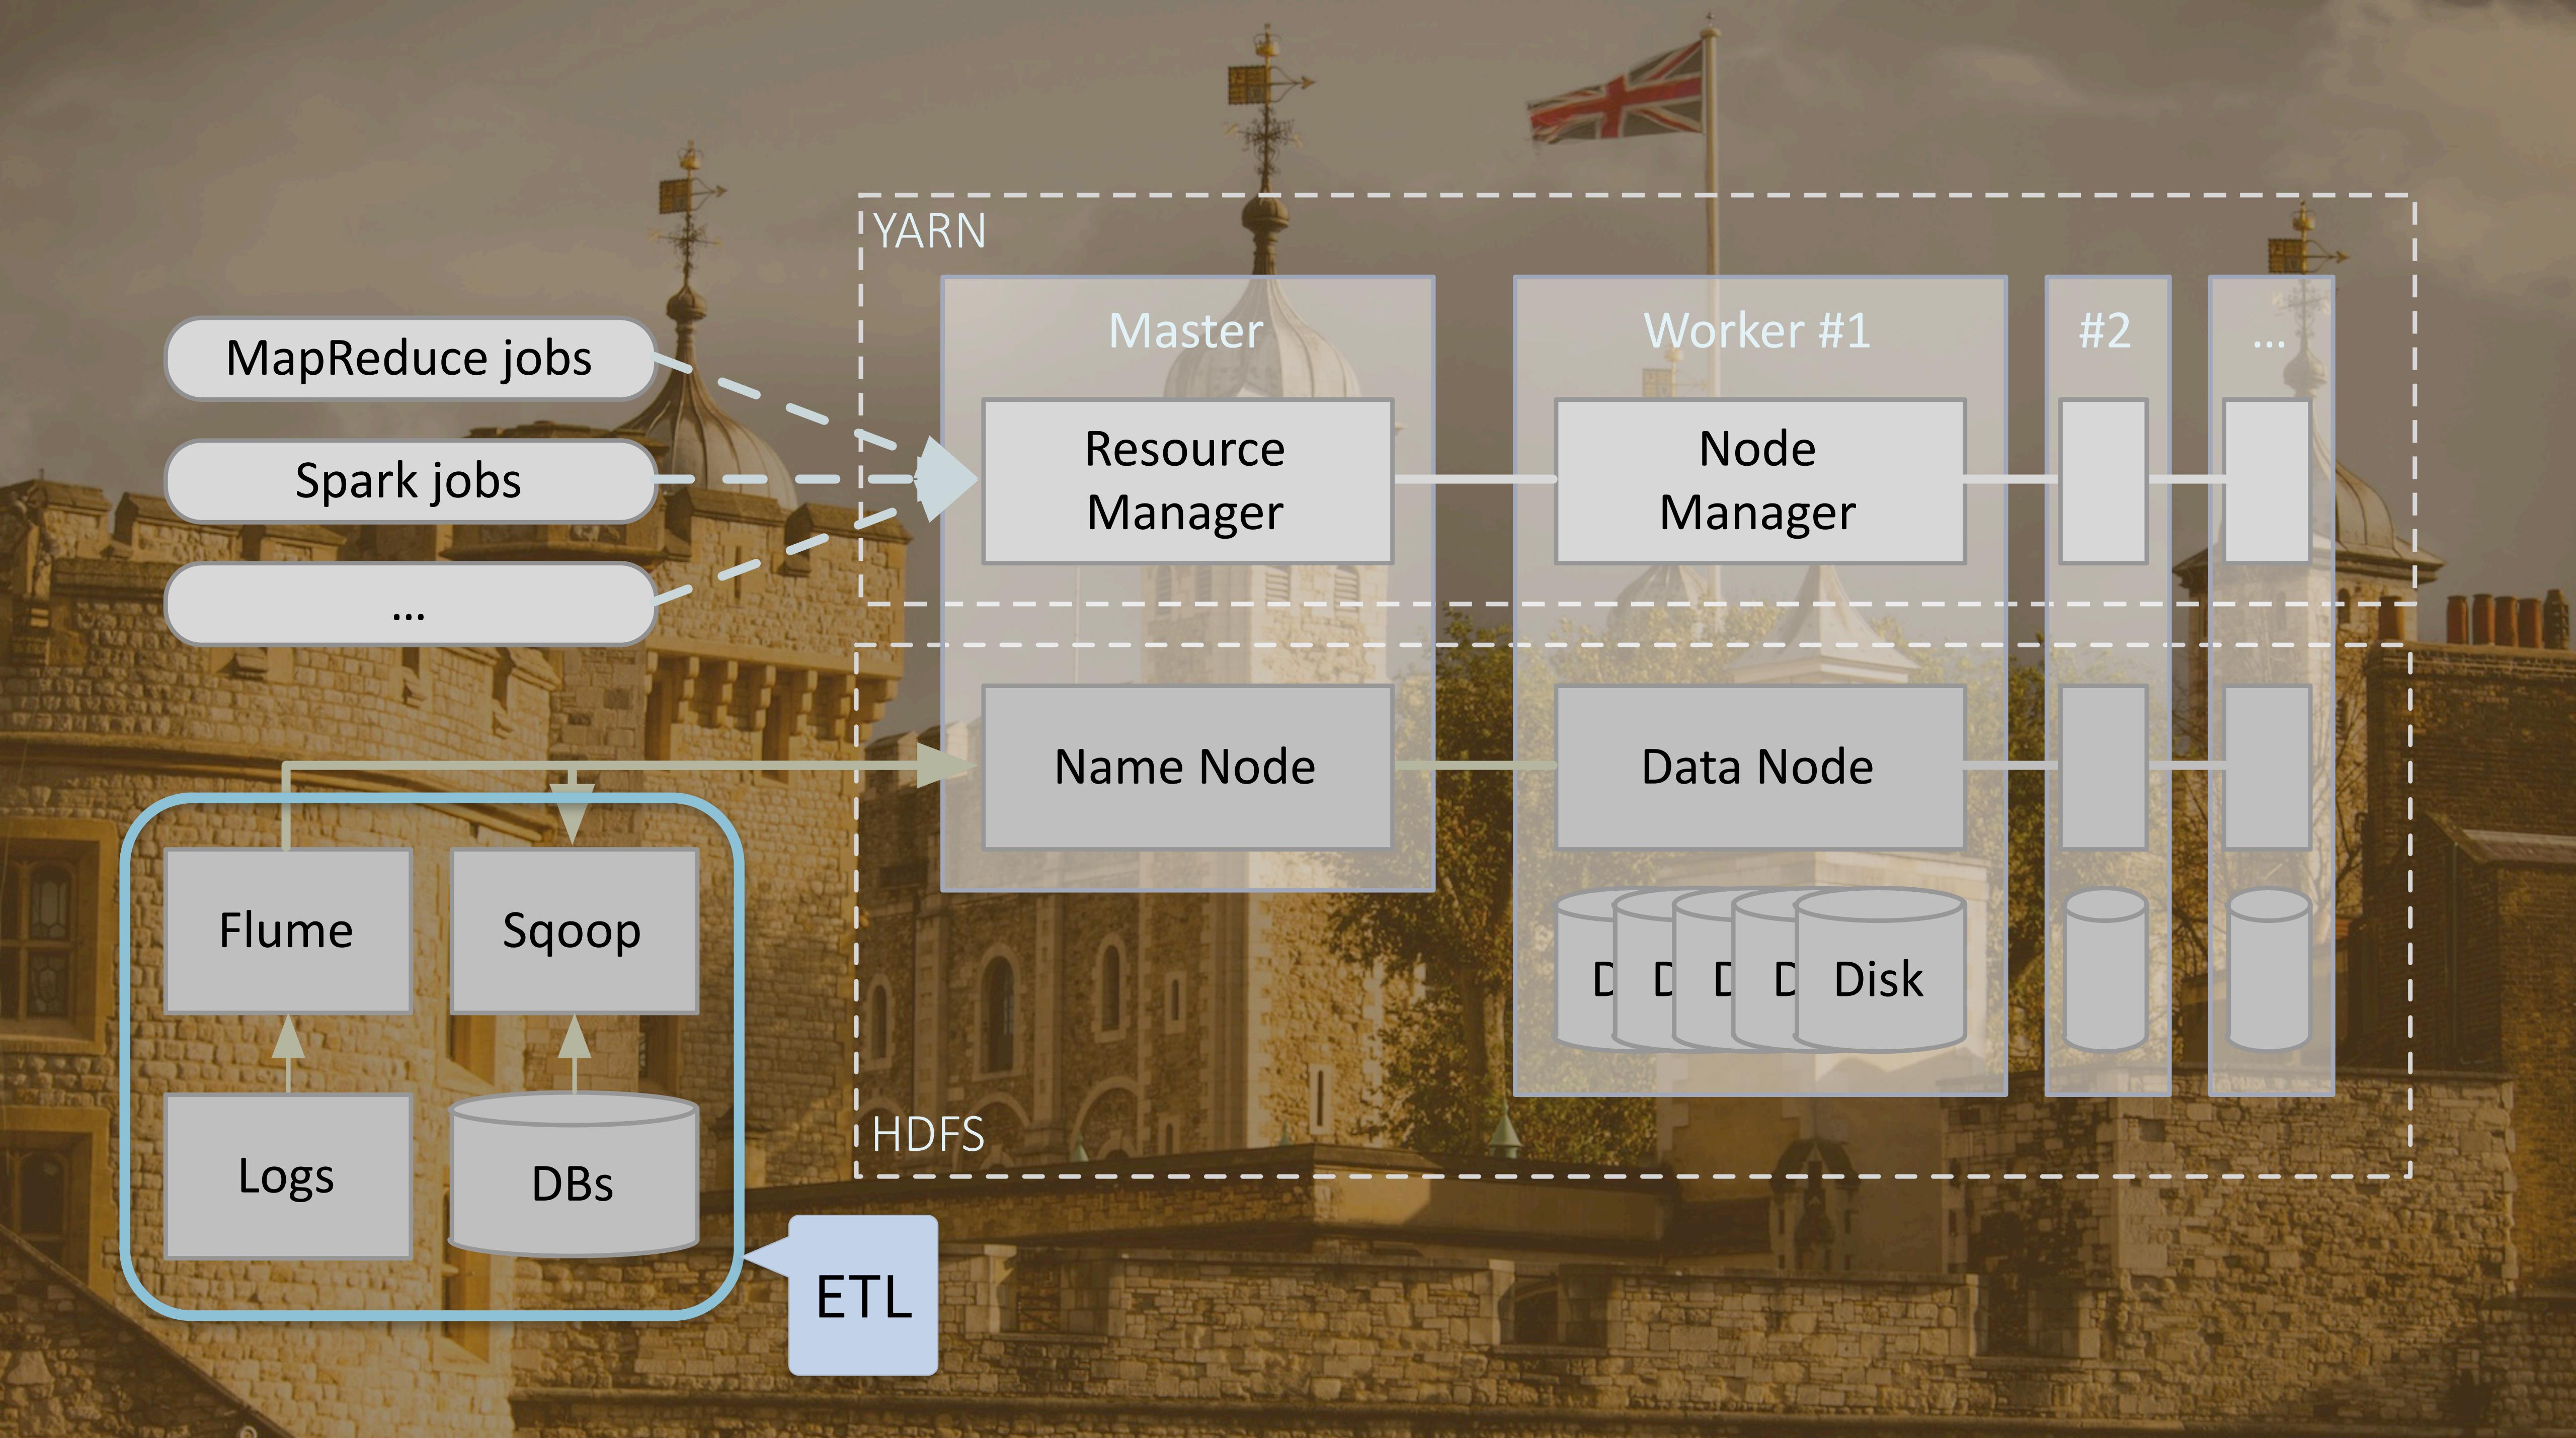


Resource Management



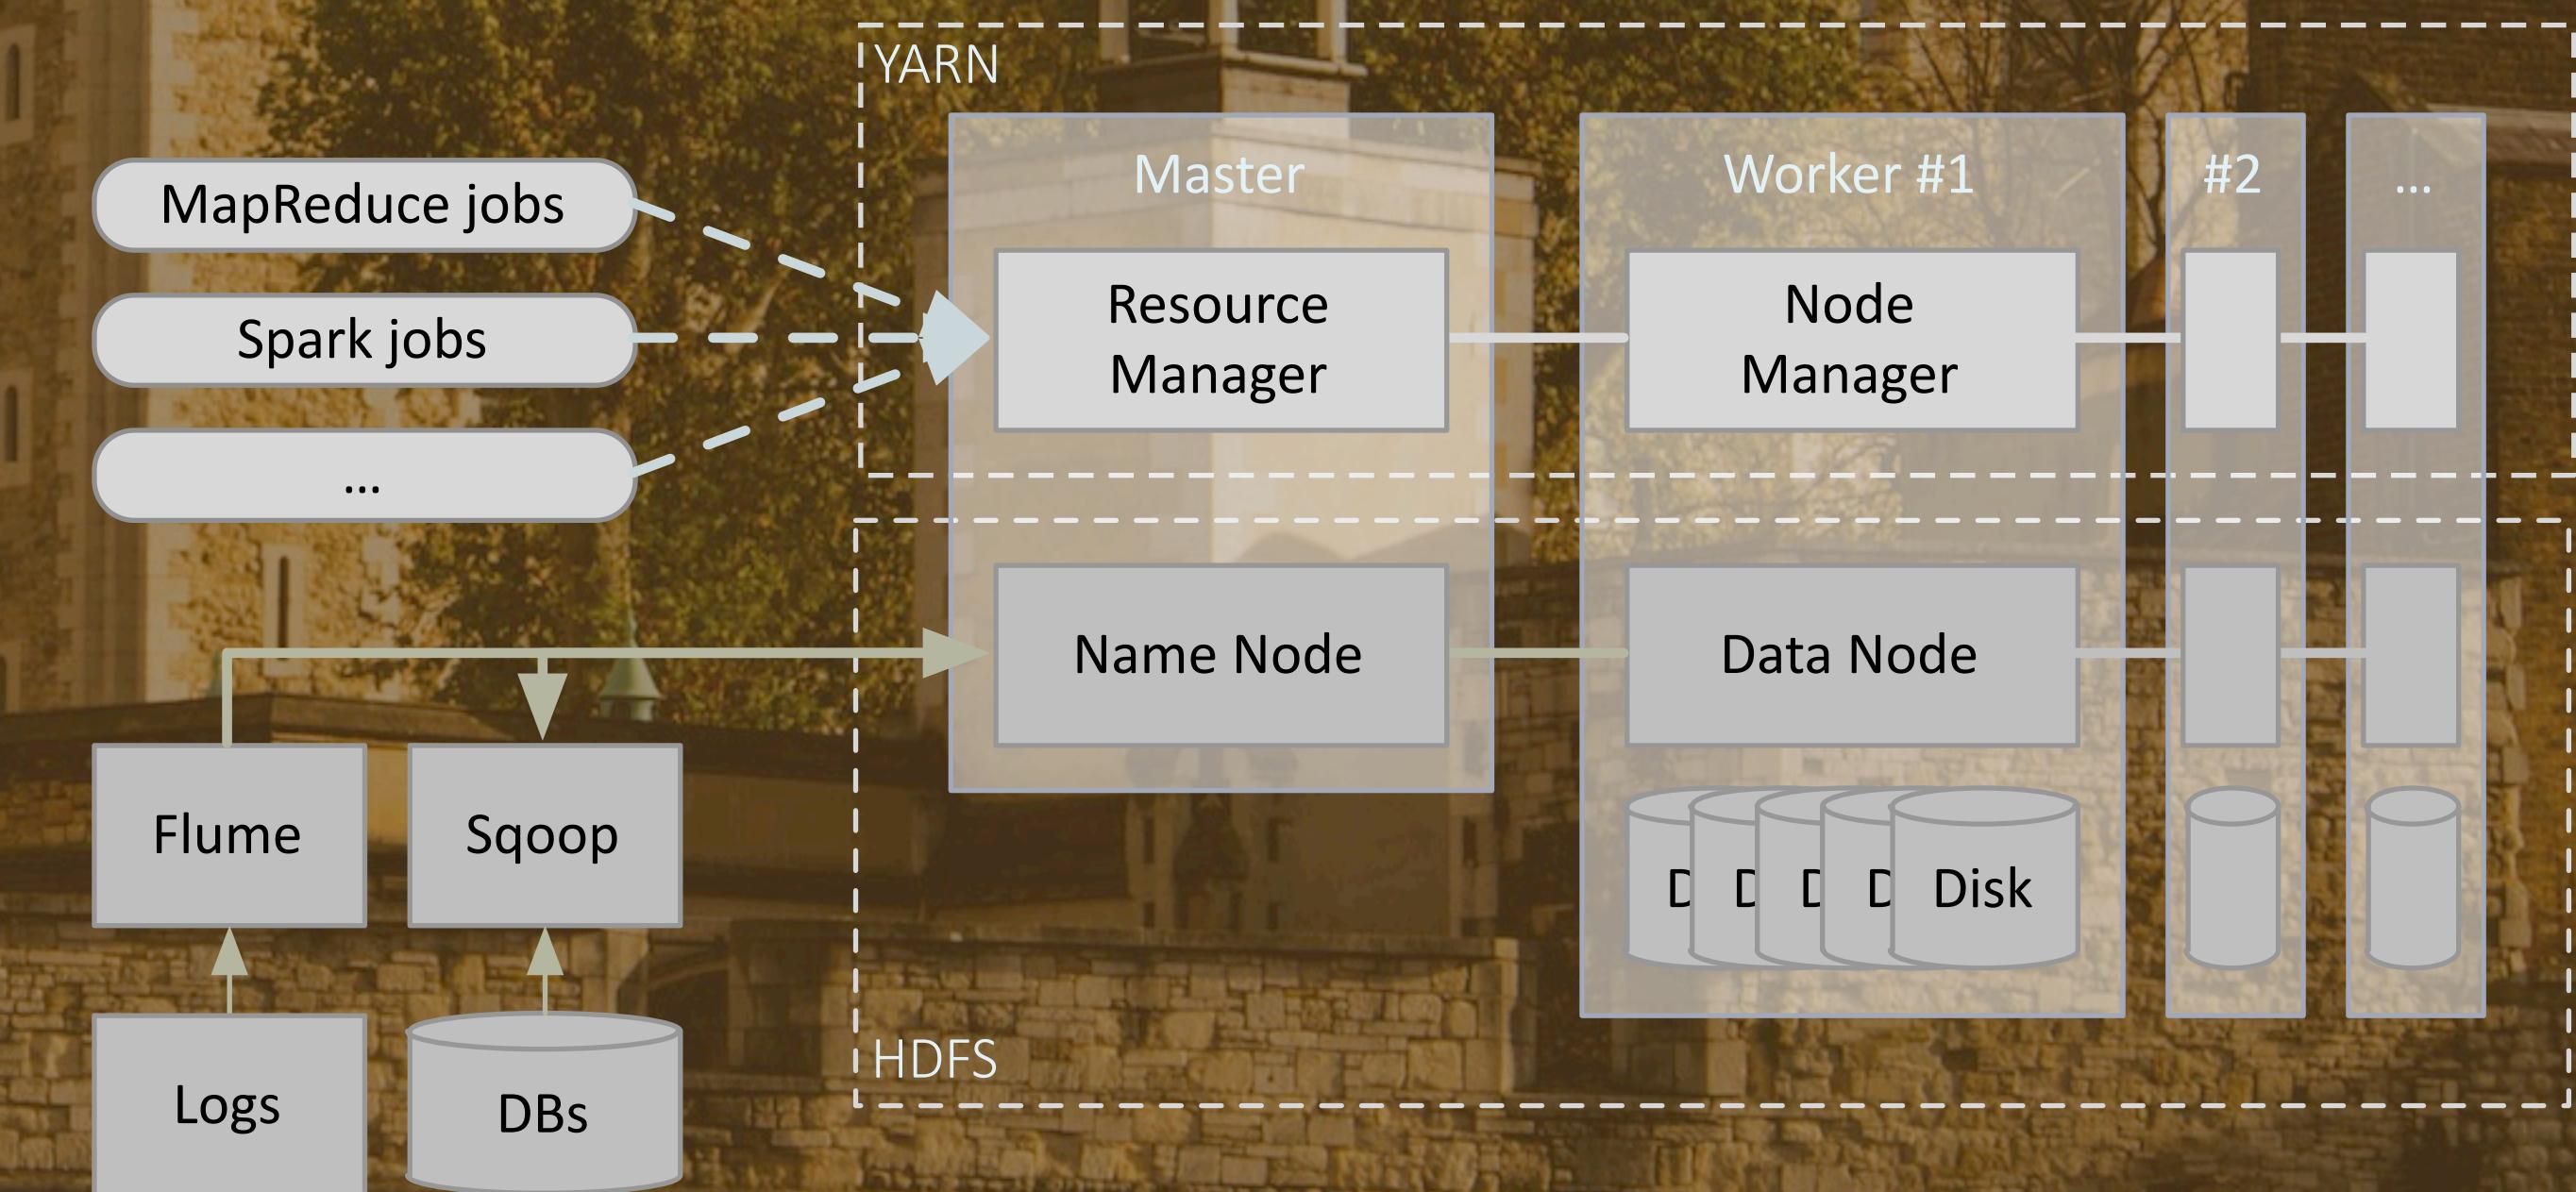
Database Deconstructed!





Other tools are built on this foundation or supplement it, like tools for ingesting data, such as Sqoop and Flume.

- 10000 meter view:
- Hadoop is the database deconstructed and reimagined
- ... but also constrained by that model

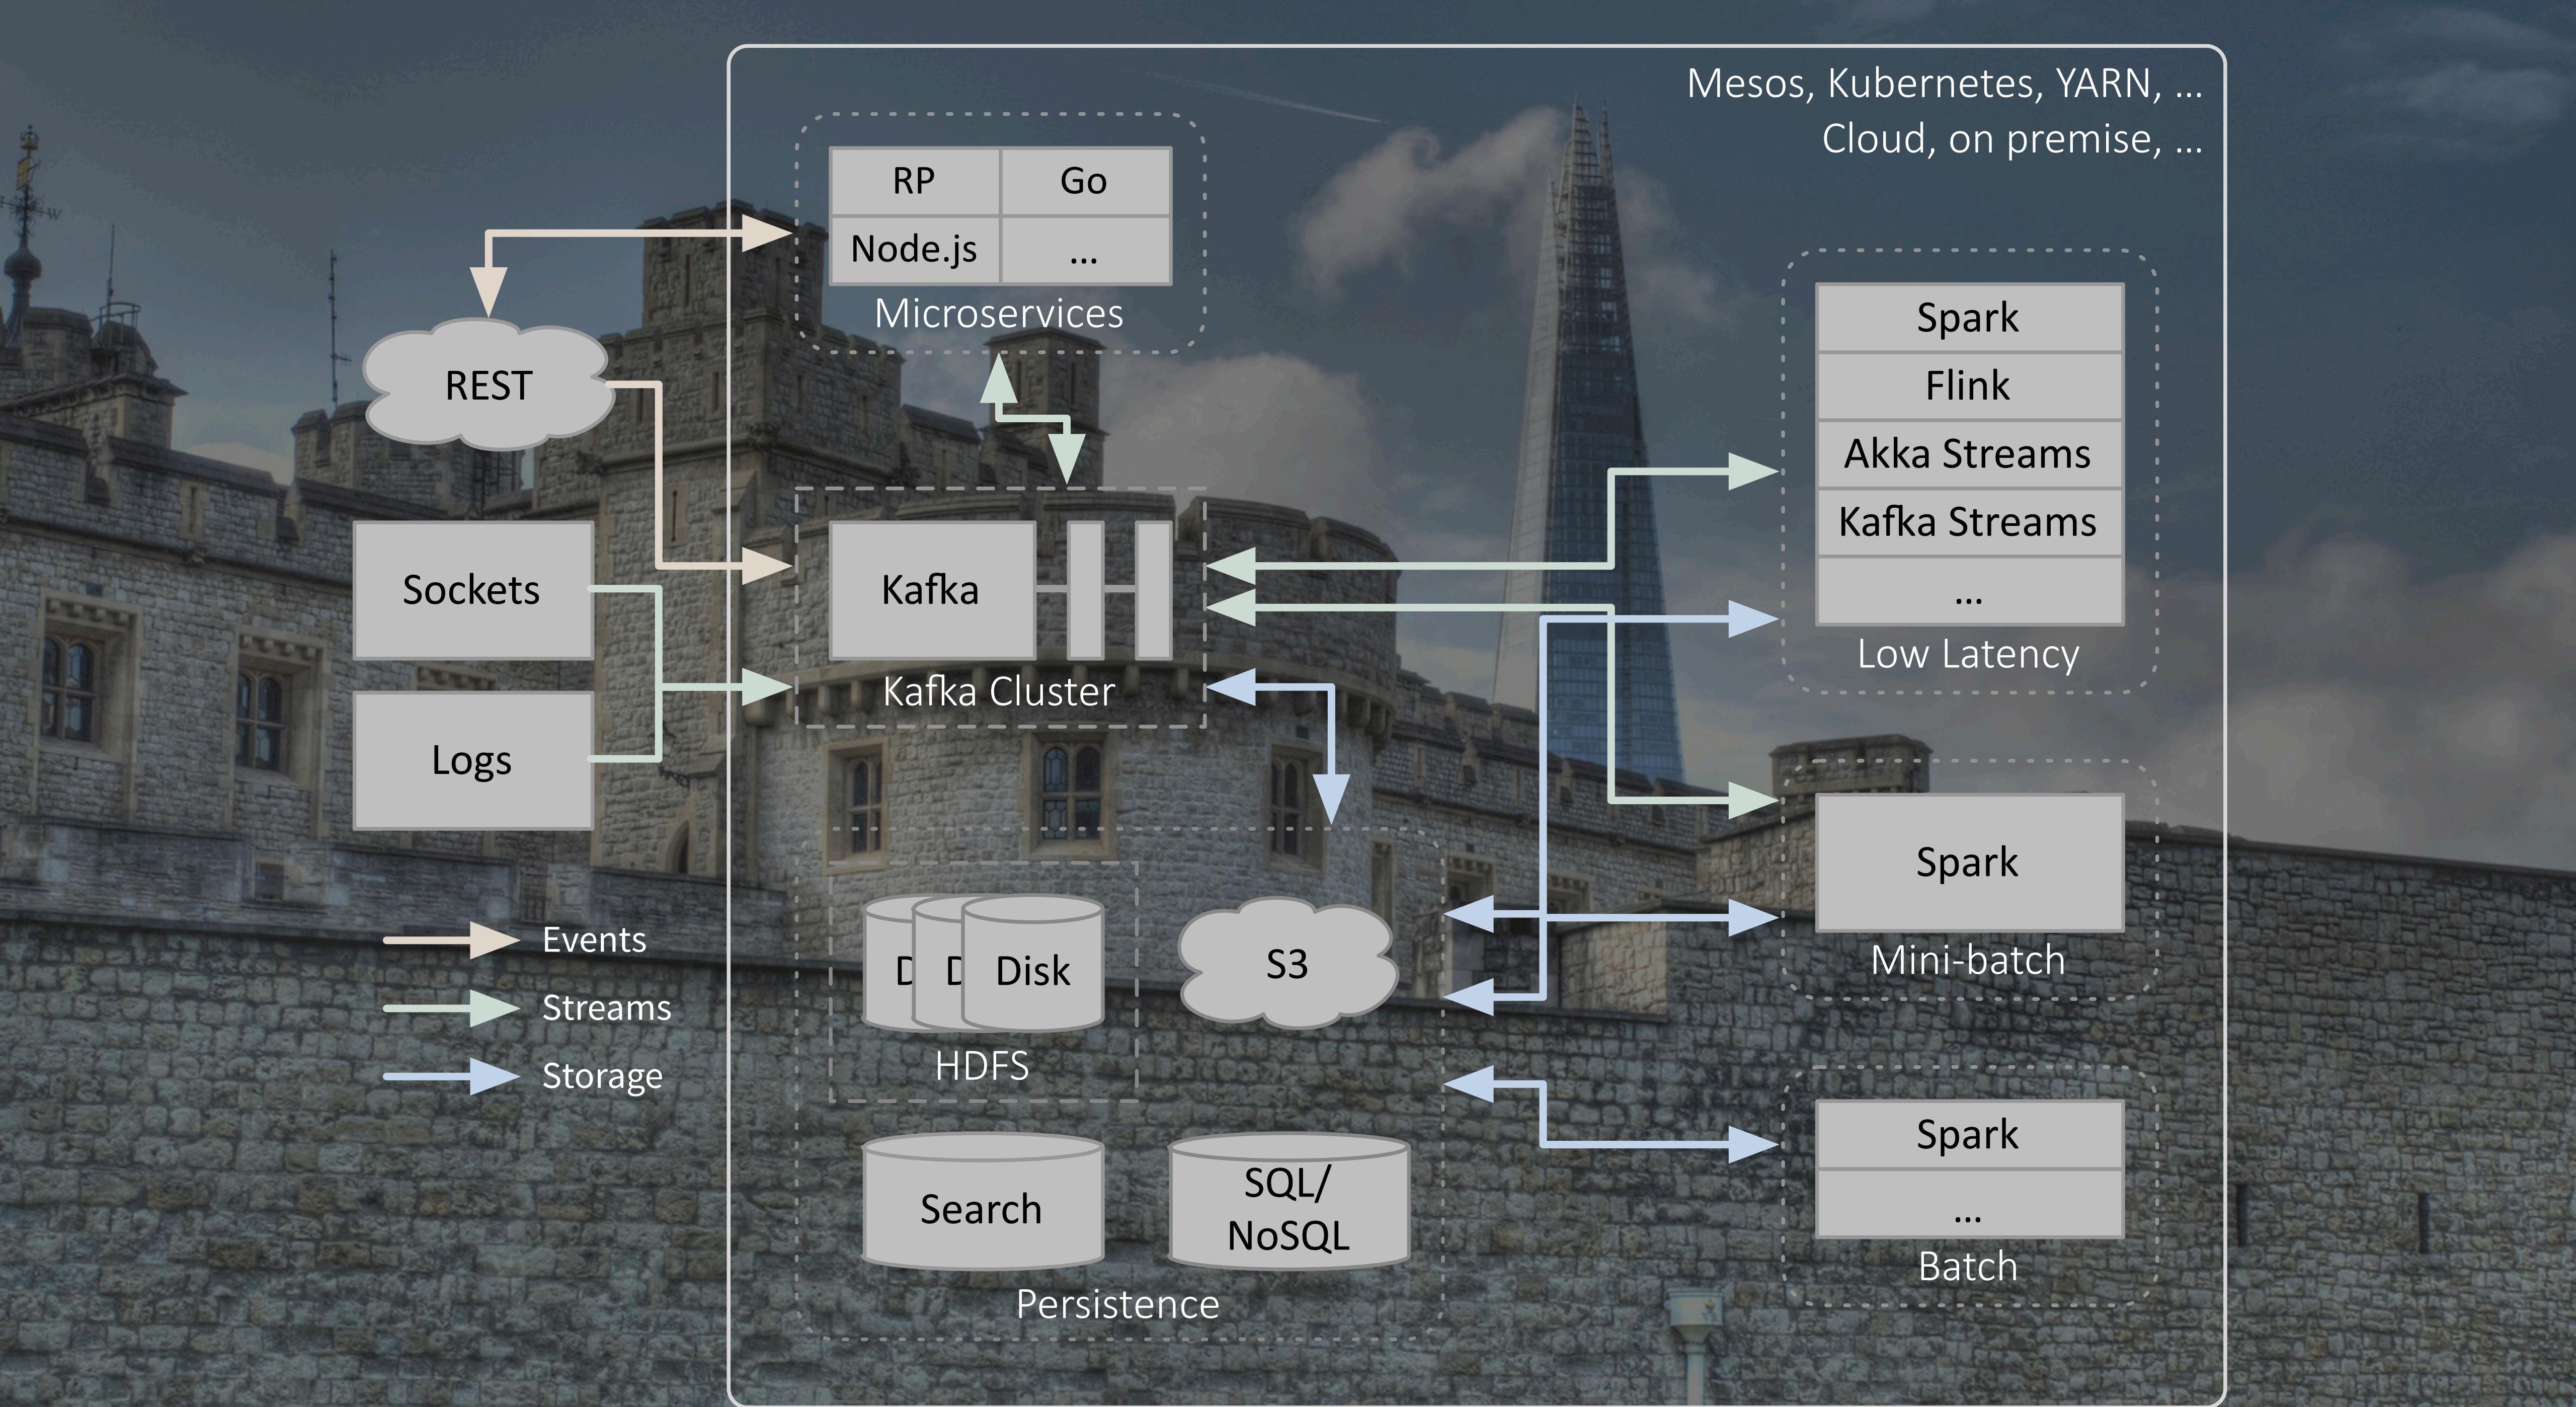


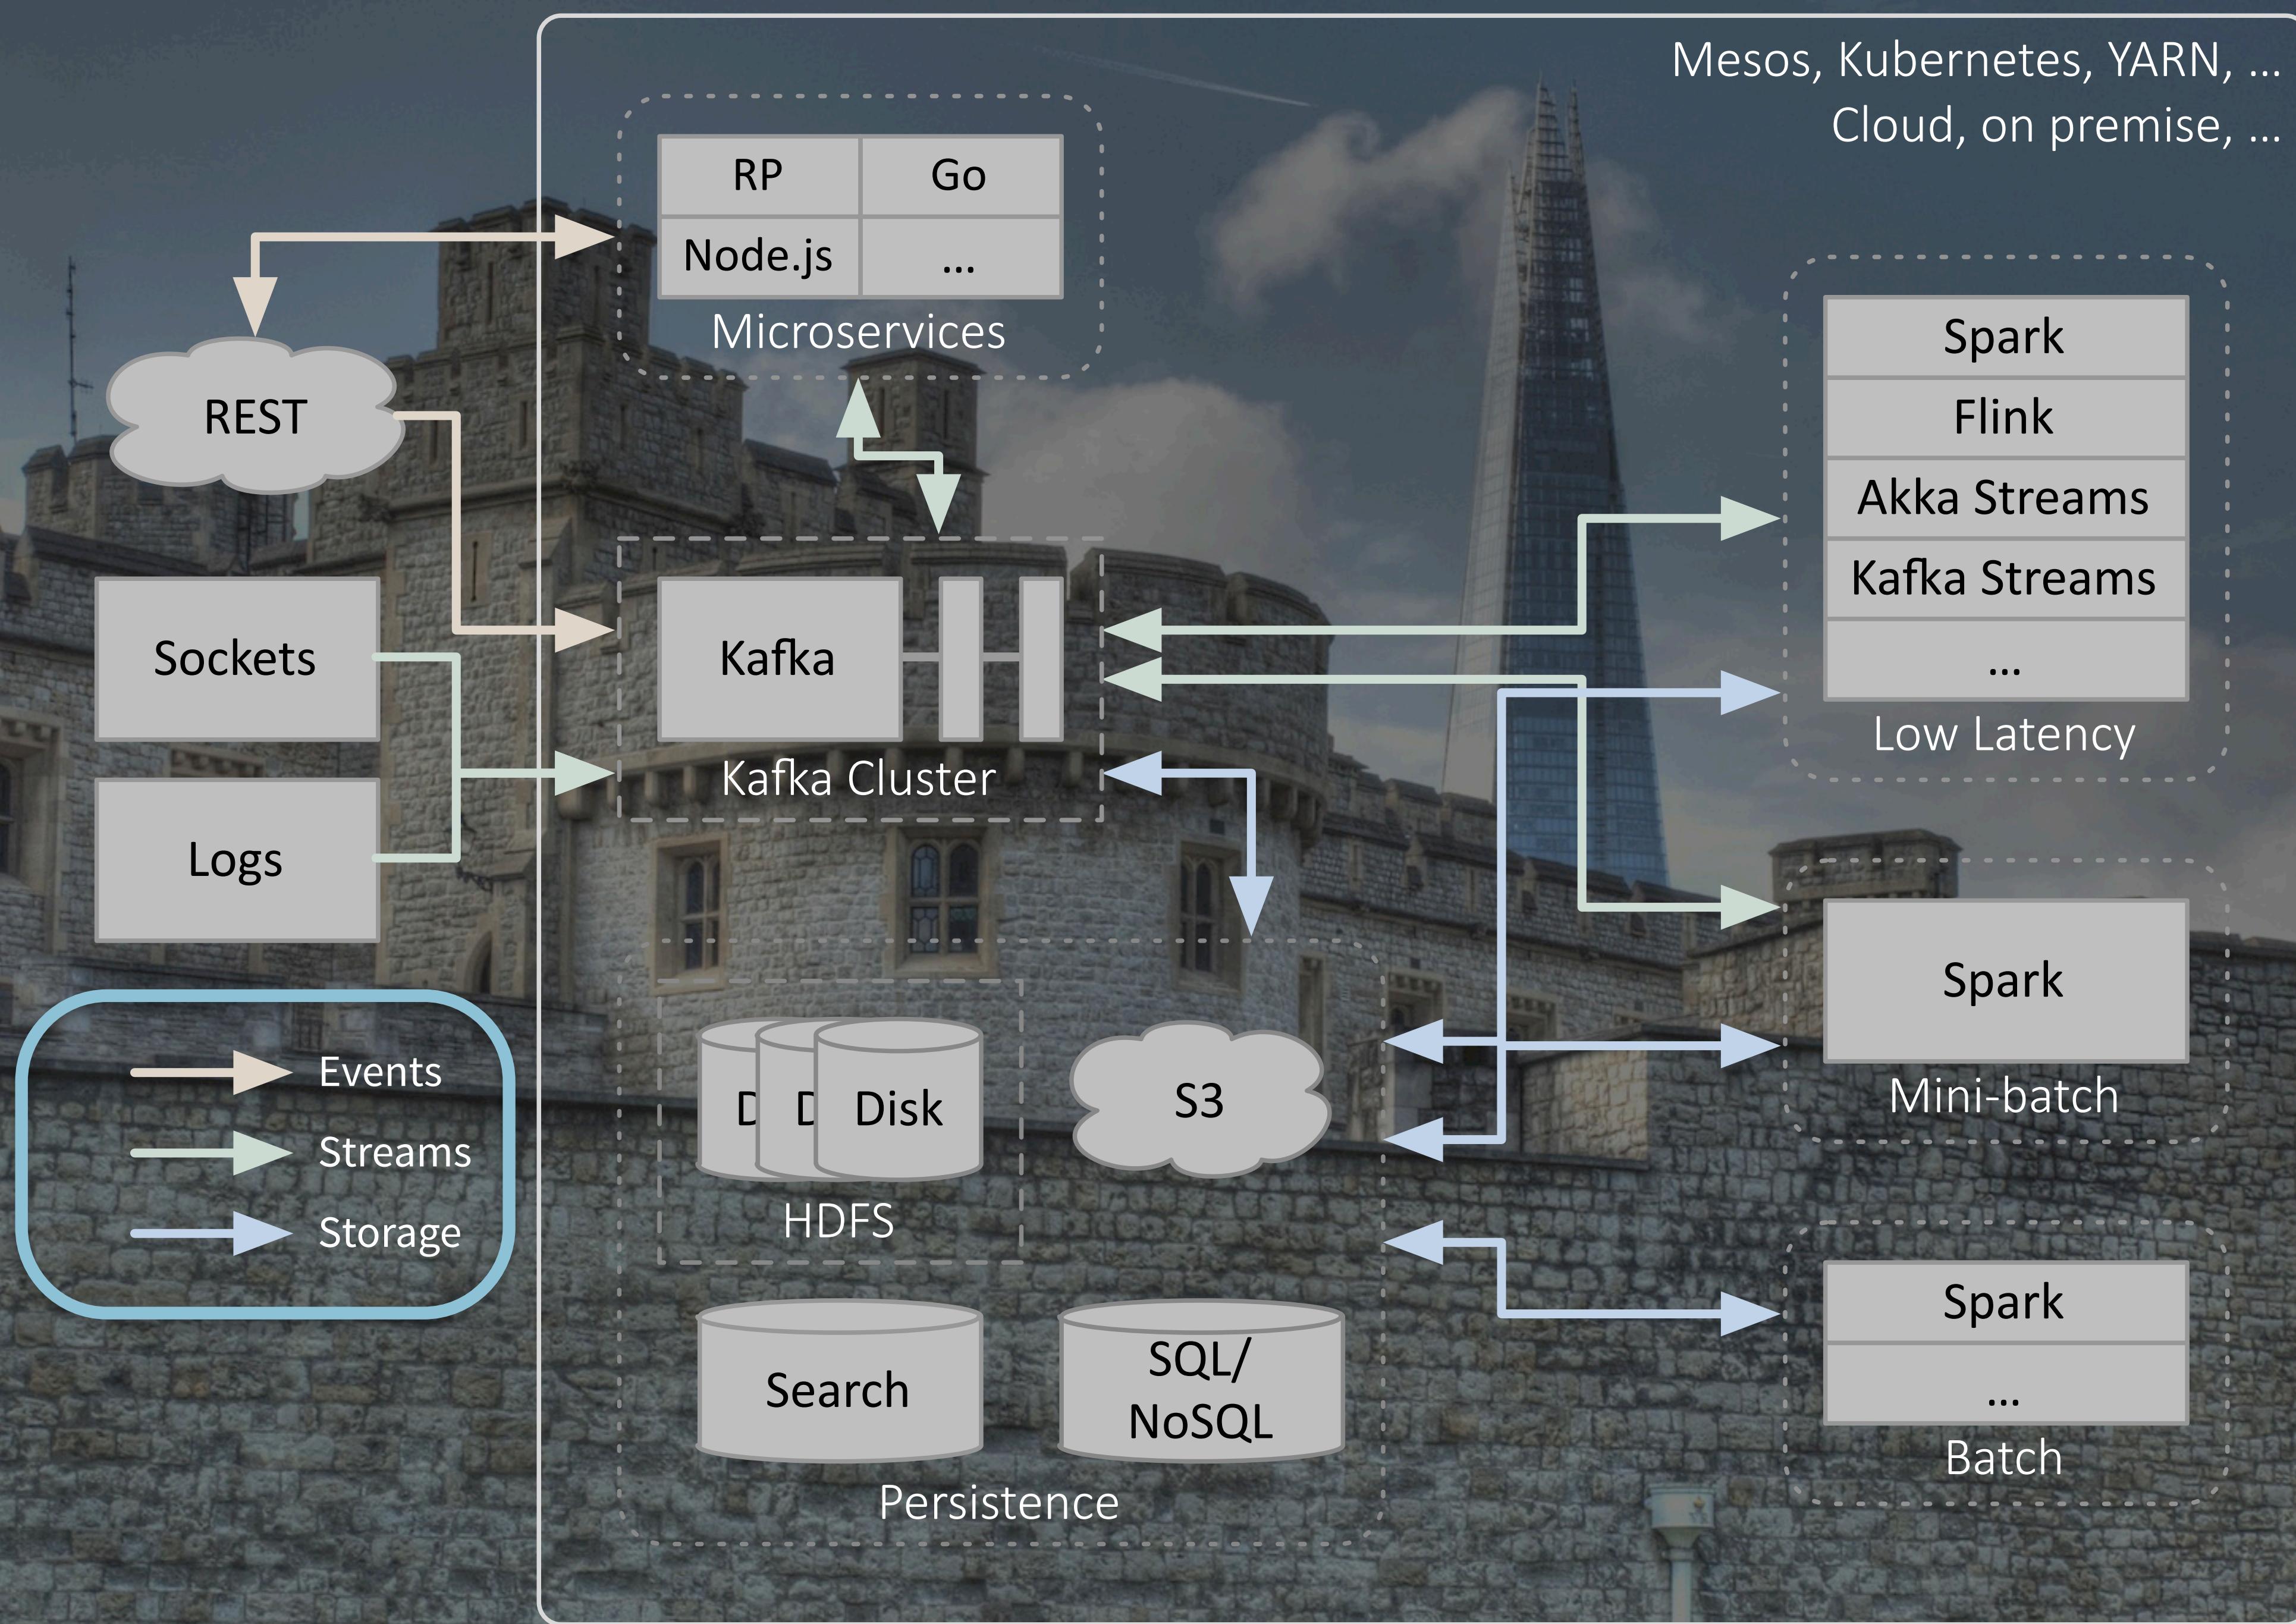
New Fast Data Architecture



But it will also support batch!!

Photo: New vs. old architectures; The Shard looming over The Tower



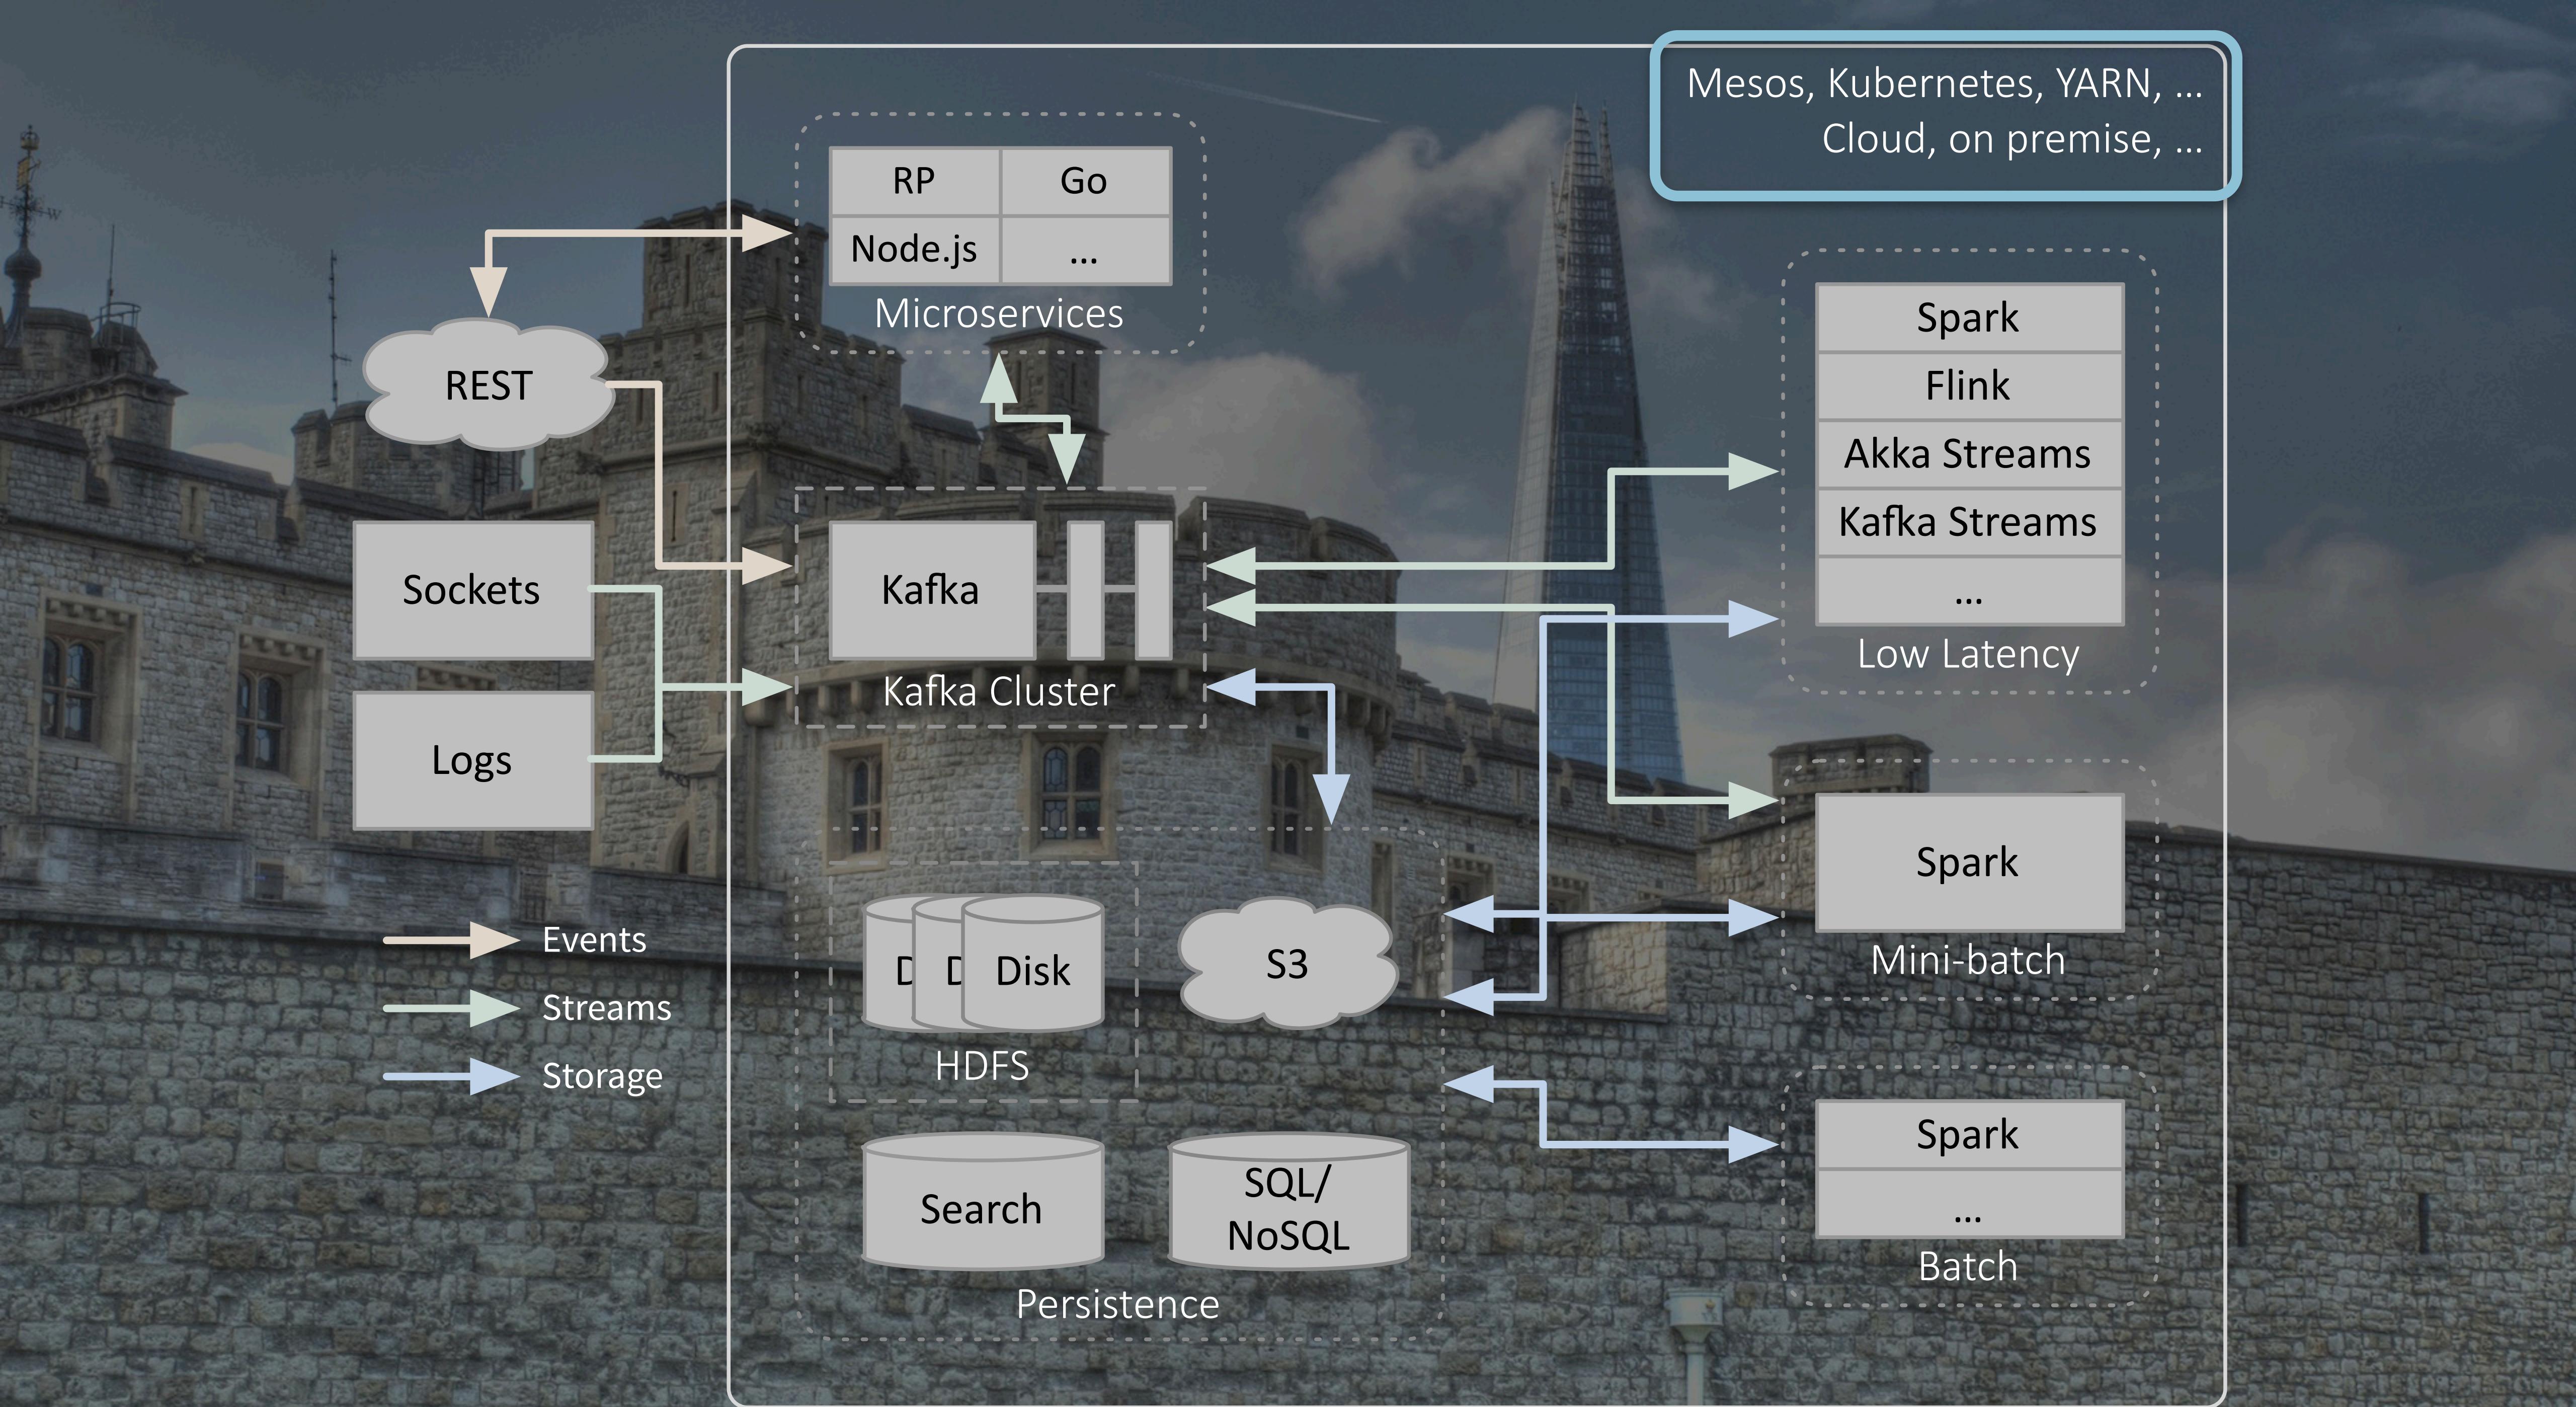


"Events" here refers to data associated with the kind of messages you use in REST, where each datum is handled uniquely (e.g., "add to cart", "purchase cart", etc.)

"Streams" could contain events or "anonymous" records, i.e., structured data like log entries, but where they are typically handled in bulk. Also, there is usually not two-way communication, except when some sort of transactional system is used to achieve "exactly once" handling, unlike REST.

"Storage" encapsulates the many ways that data is written to and read from durable stores.

The differences are mostly intended to capture different volumes, latency characteristics, communication protocols, etc.



You can run this architecture on Mesos, Kubernetes, even YARN (Hadoop), although YARN is not really flexible enough to run the wide variety of services in this picture. You can also deploy on premise or in cloud environments. At Lightbend, we have a lot of customers who are standing up these fast data clusters beside their Hadoop clusters, reserving the latter for the batch and interactive workloads described previously.

Send events,
manage sessions,
less about “fast
data”

REST

Sockets

Logs

Fire and forget
streams of data;
capture it or you
lose it (e.g.,
Twitter firehose)

Events
Streams
Storage

Structured data
like server logs,
click-stream logs,
etc.

Mesos, Kubernetes, YARN, ...
Cloud, on premise, ...

Spark

Flink

Akka Streams

Kafka Streams

...

Low Latency

Spark

Mini-batch

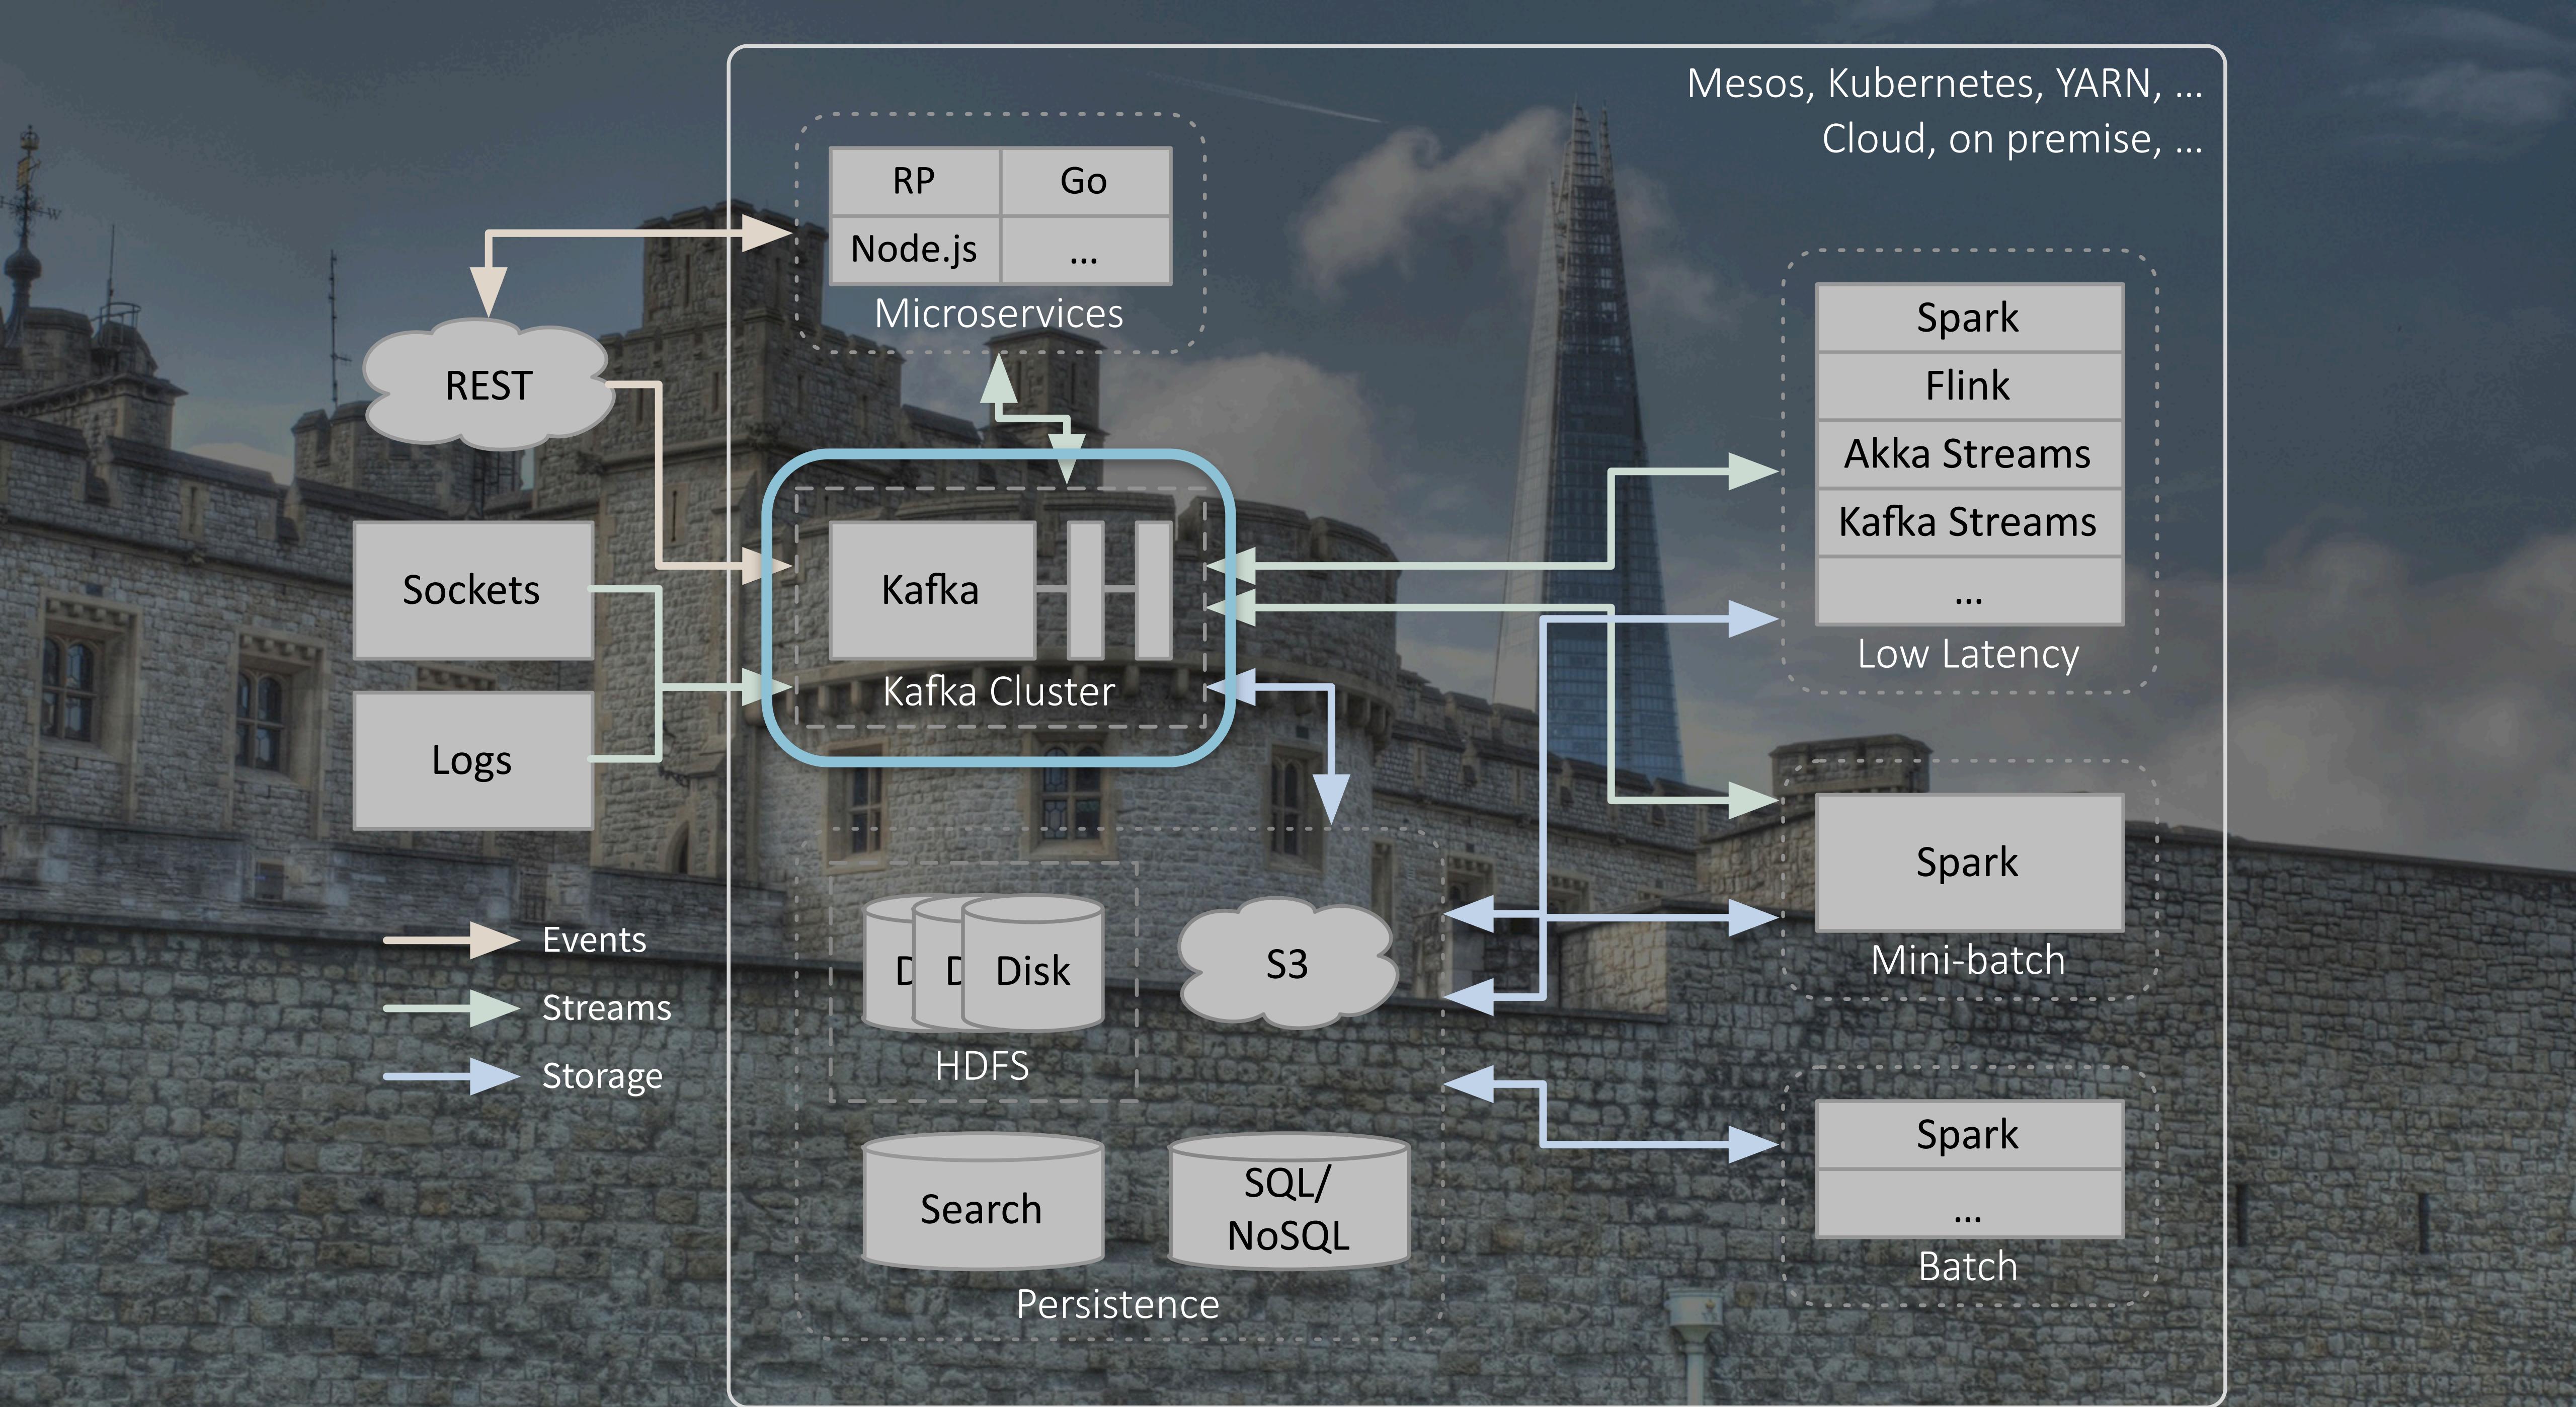
Spark

...

Batch

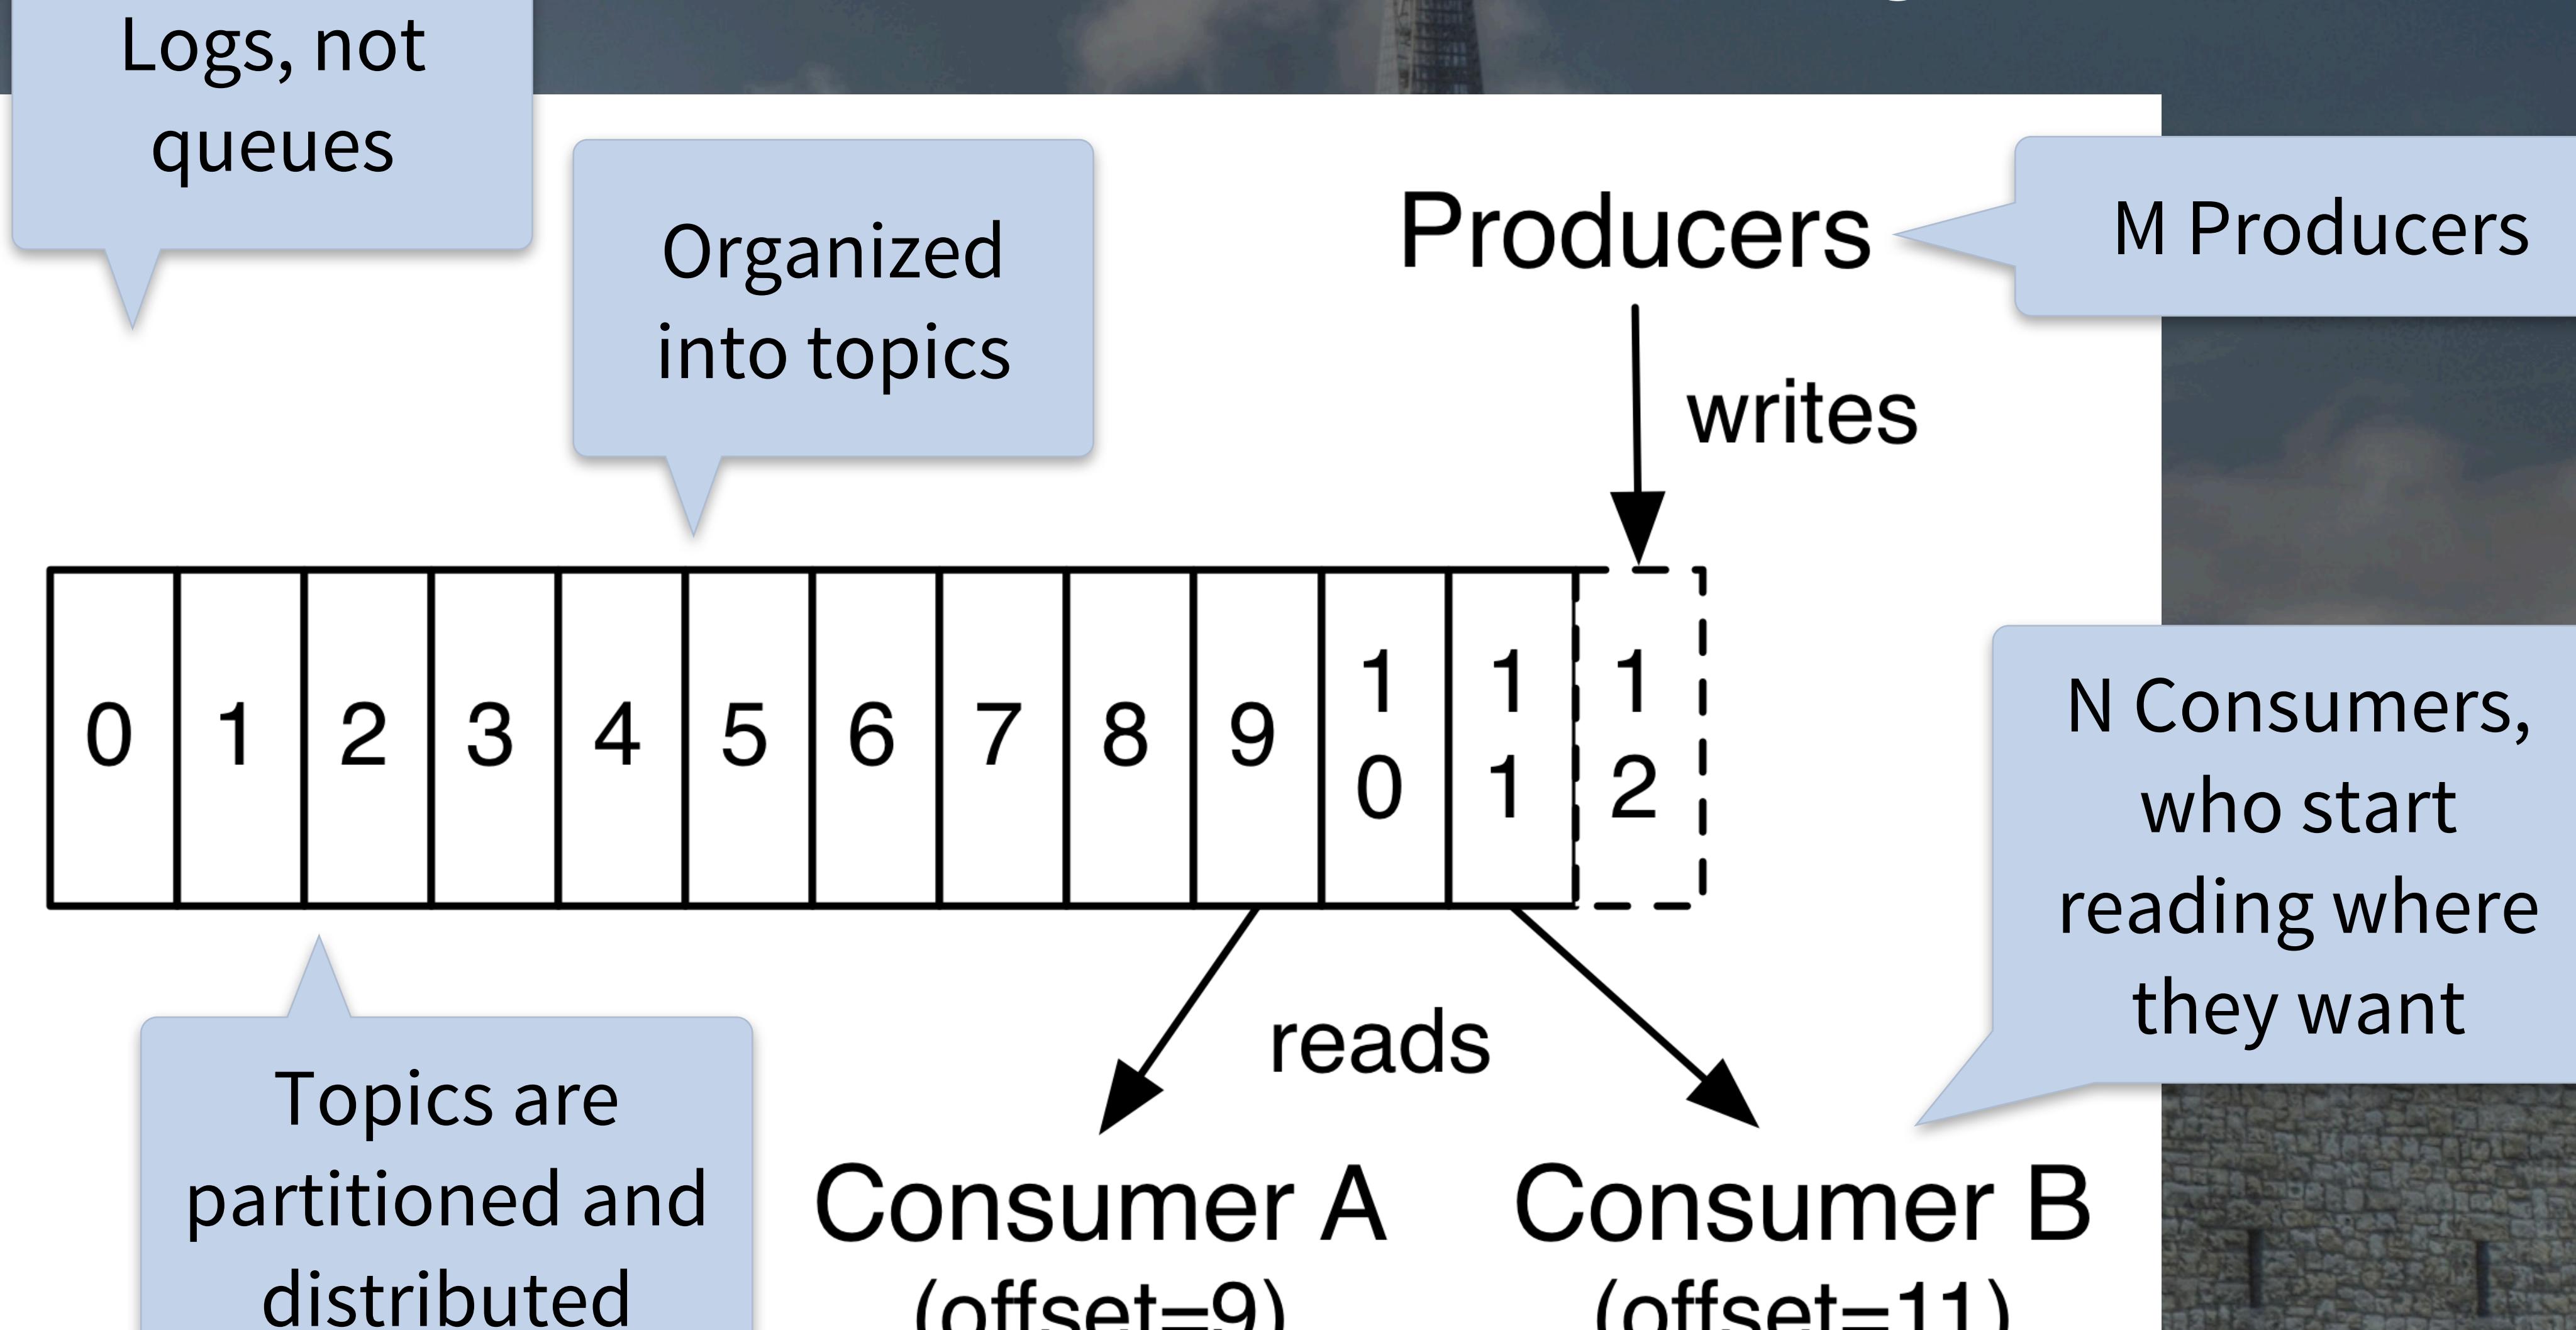
S3

SQL/
NoSQL



Kafka is the core “backplane” of this architecture, the place where data is ingested in log-oriented storage, organized into topics. They are not quite “queues”, like traditional message queue environments, because each consumer doesn’t delete the message it reads. Rather, Kafka manages message lifecycles, so that N consumers can see the whole stream. M Publishers and N consumers are decoupled for any numbers M and N. Kafka has massive scalability and excellent resiliency and data durability. All services can communicate through each other using Kafka, too, rather than having to manage arbitrary point-to-point connections.

Why Kafka?



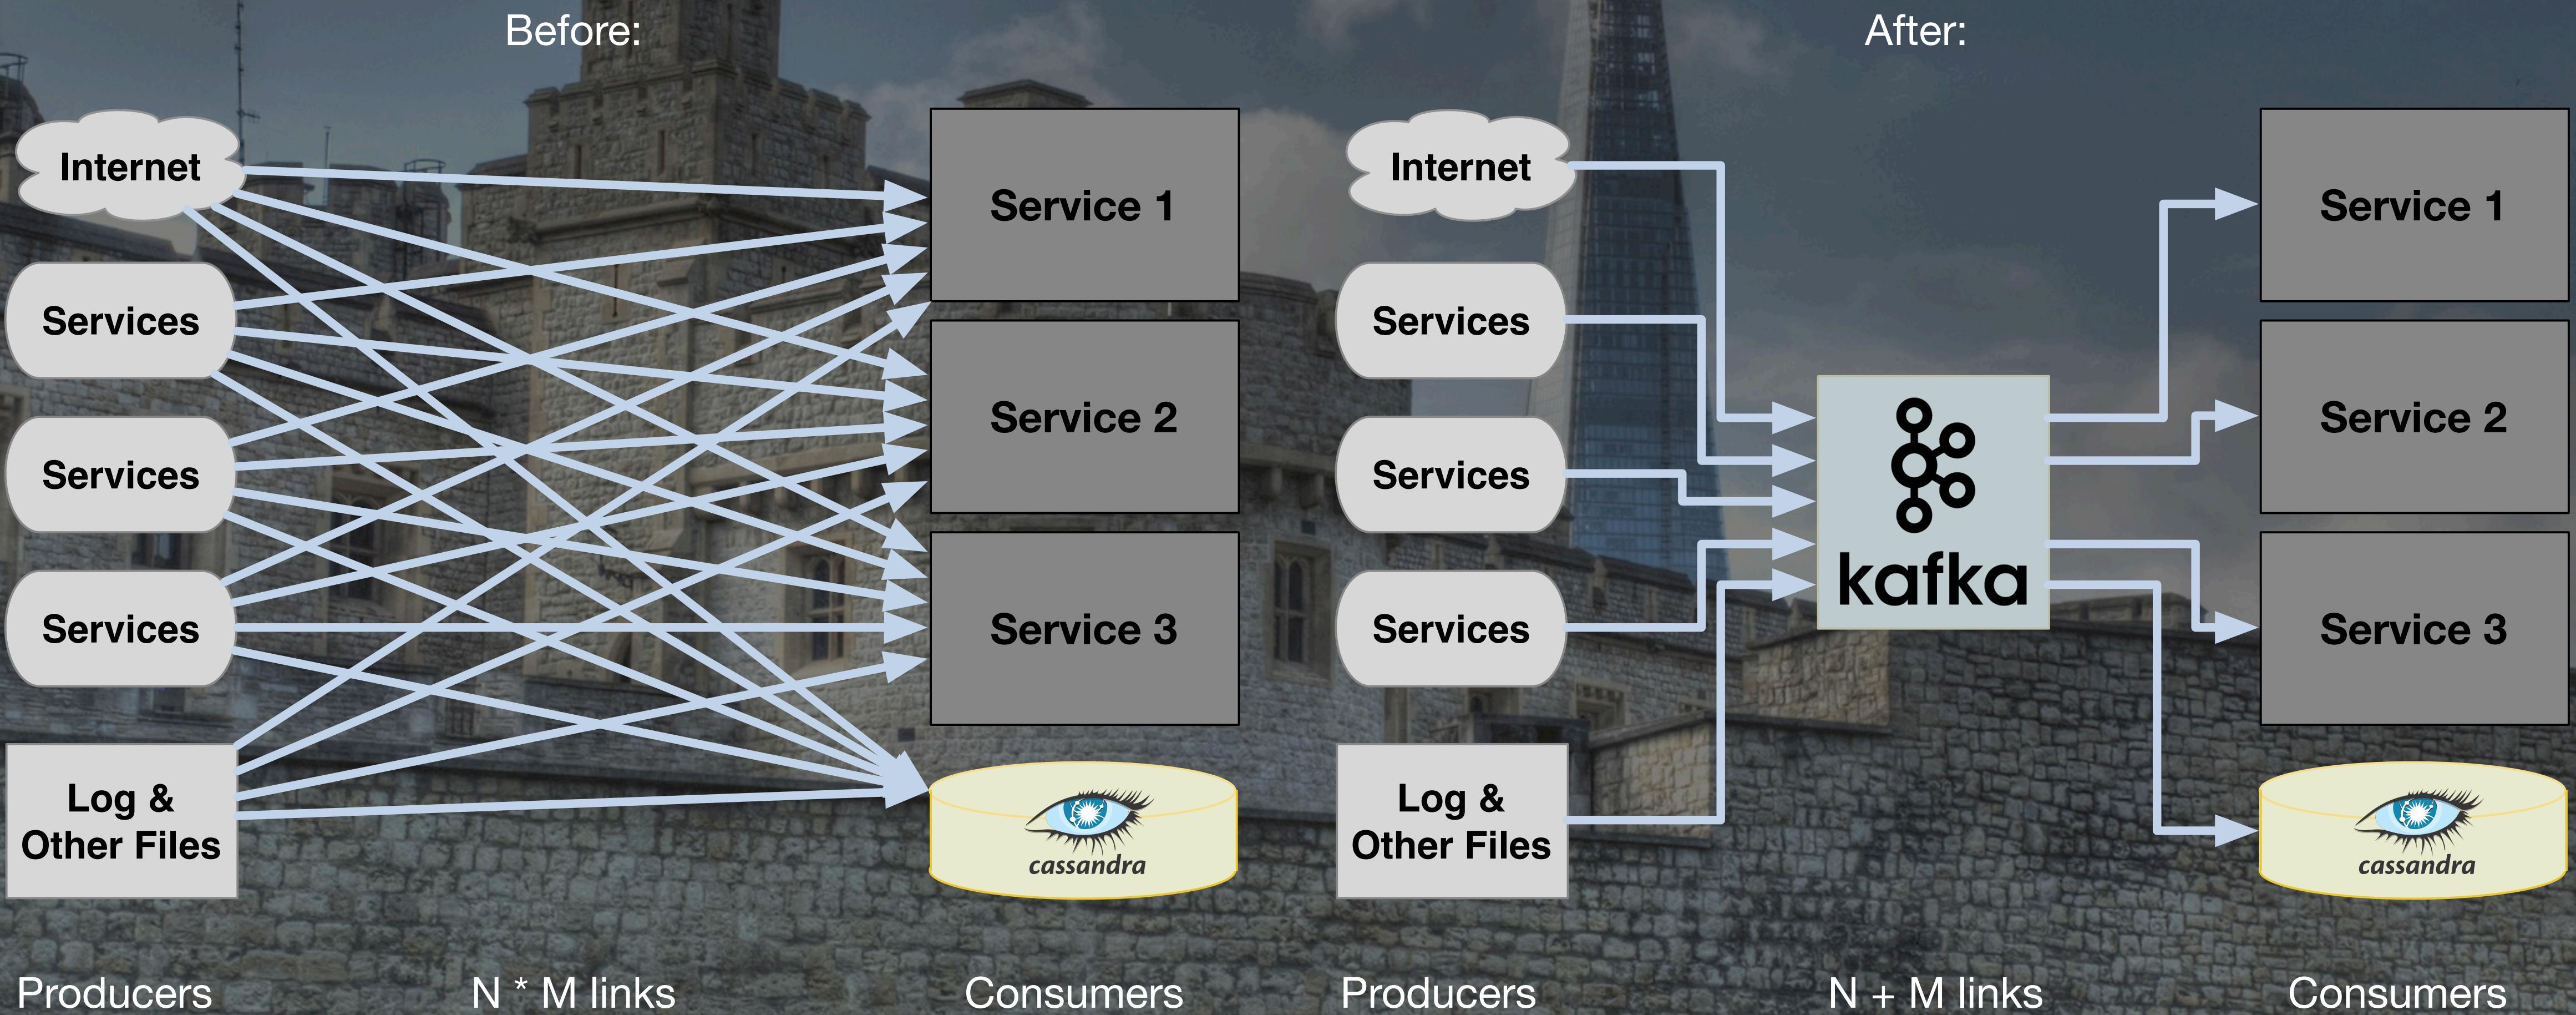
Kafka is a distributed log, storing messages sequentially. Producers always write to the end of the log, consumers can read on the log offset that they want to read from (earliest, latest, ...)
Kafka can be used as either a queue or pub sub

The main differences are:

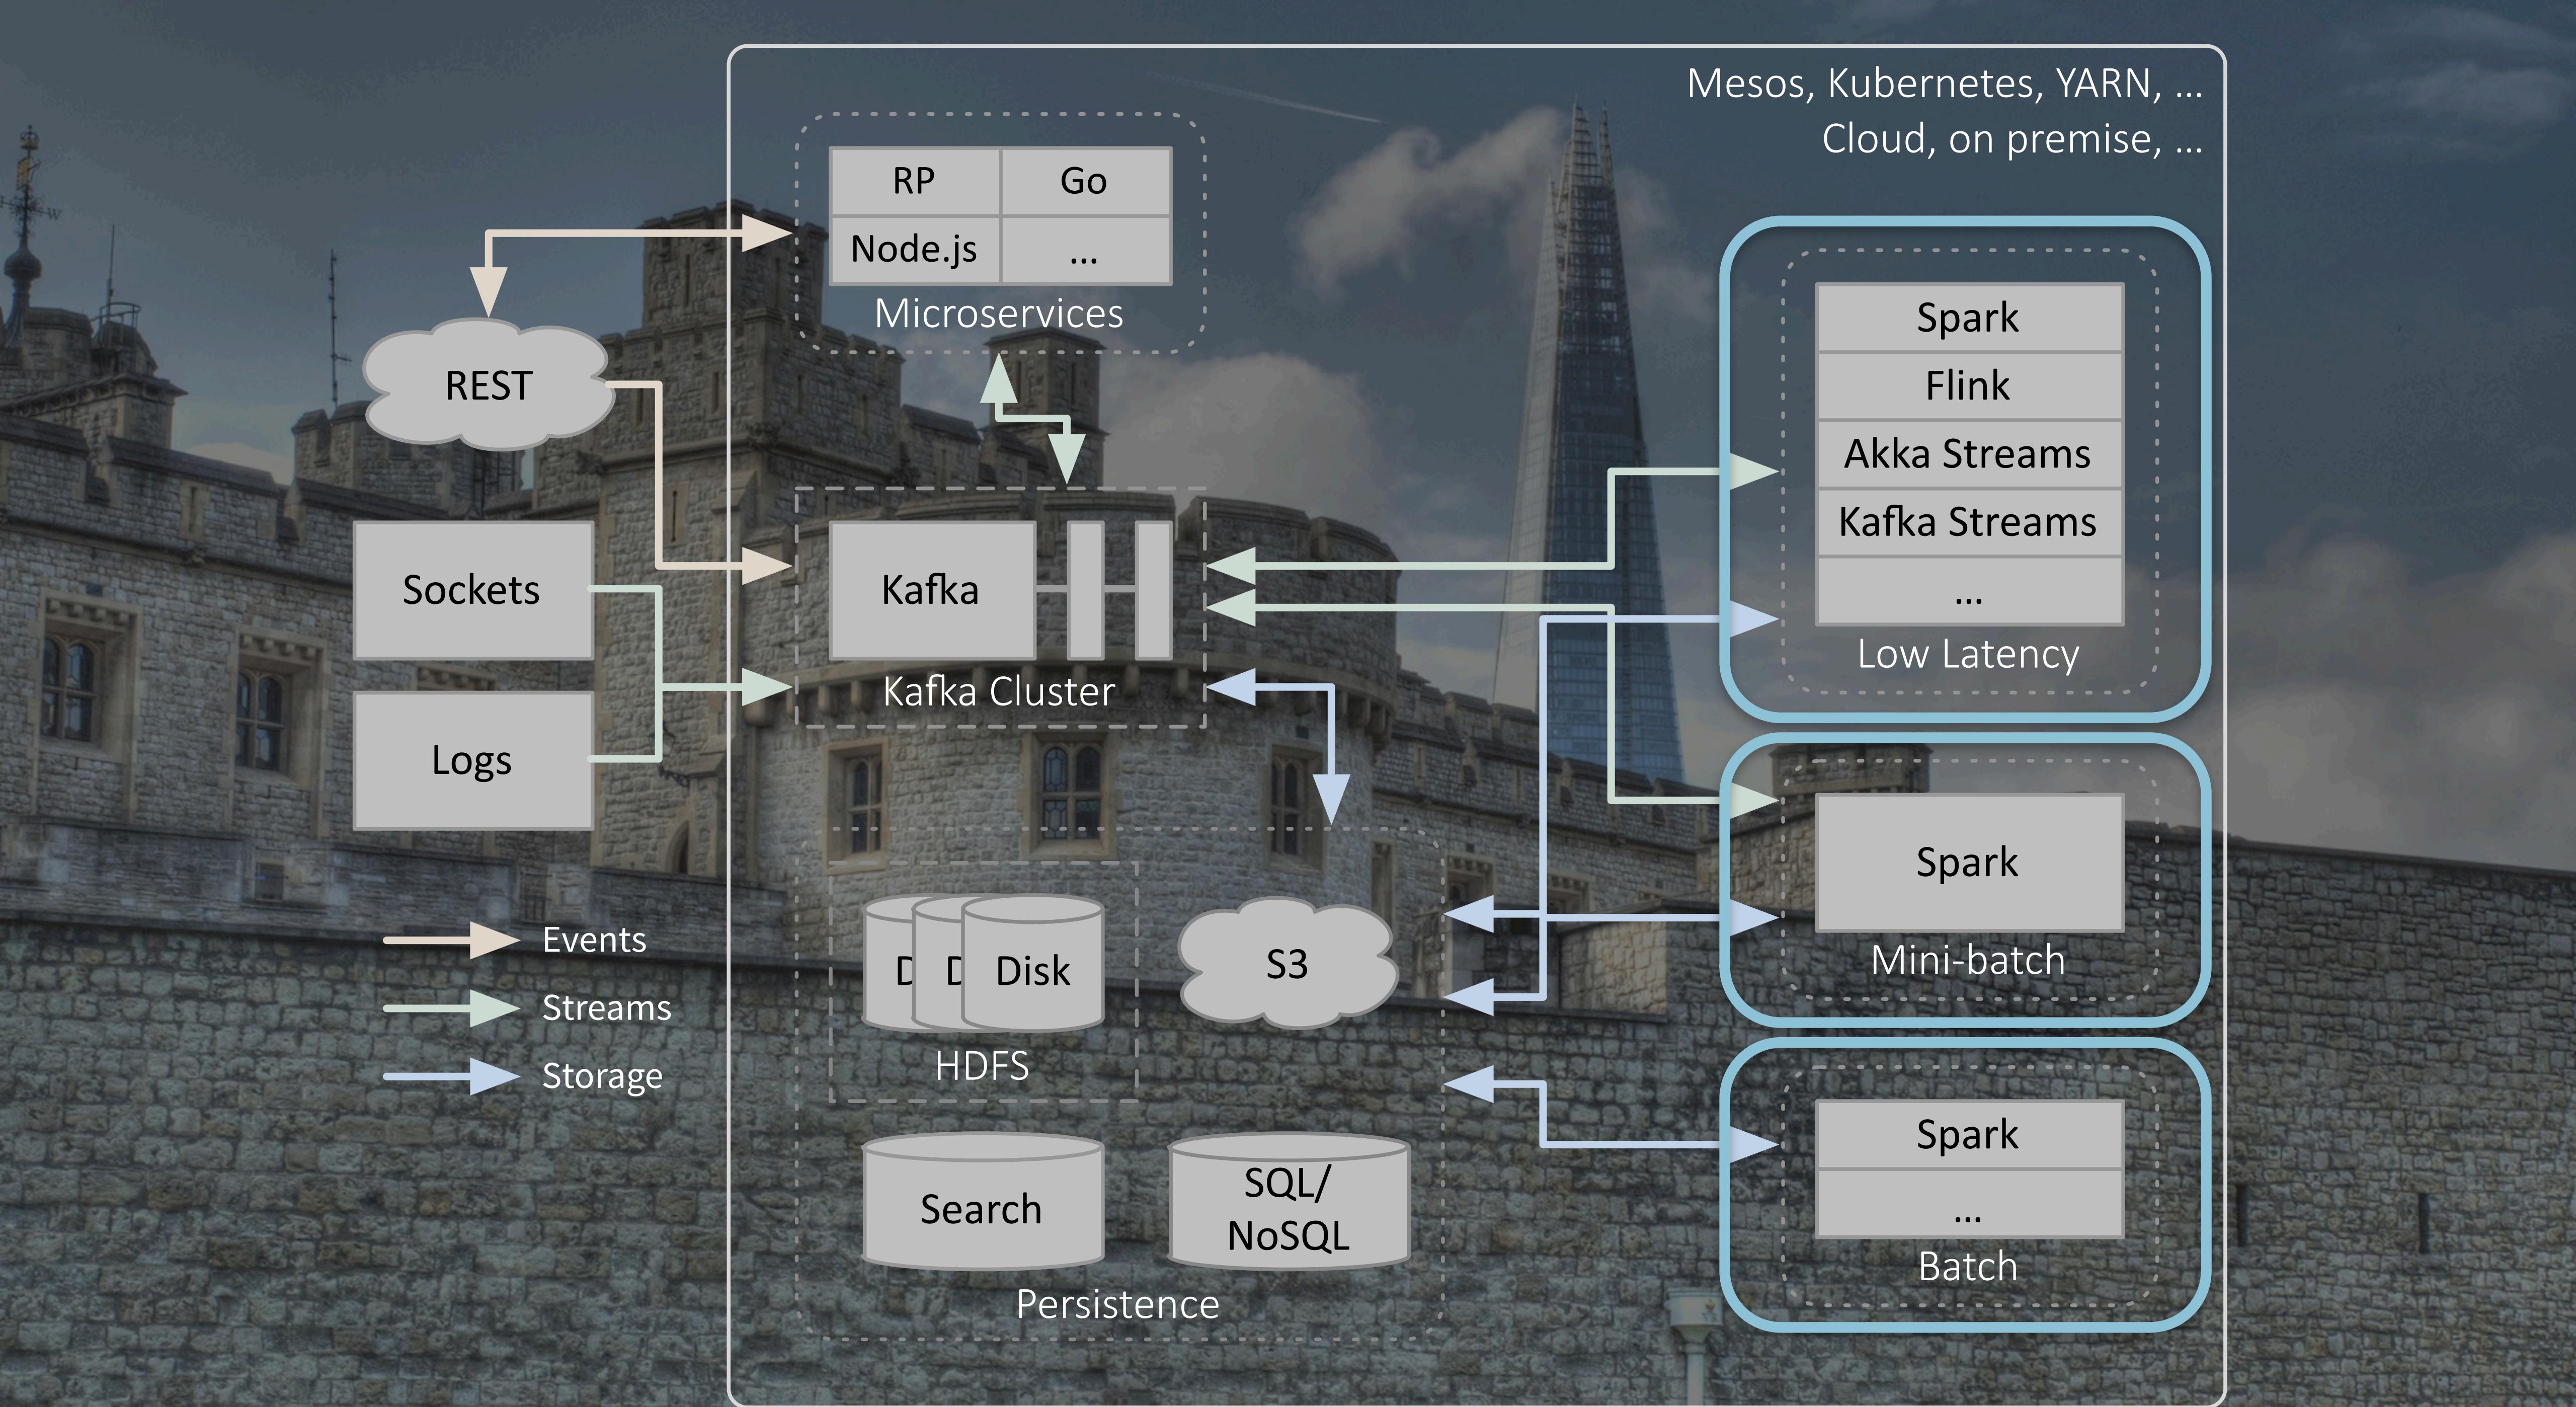
1. Log is persistent where queue is ephemeral (reads pop elements)
2. Traditional message brokers manage consumer offsets, while log systems allow users to manage offsets themselves

Alternatives to Kafka include Pravega (EMC) and Distributed Log/Pulsar (Apache)

Using Kafka



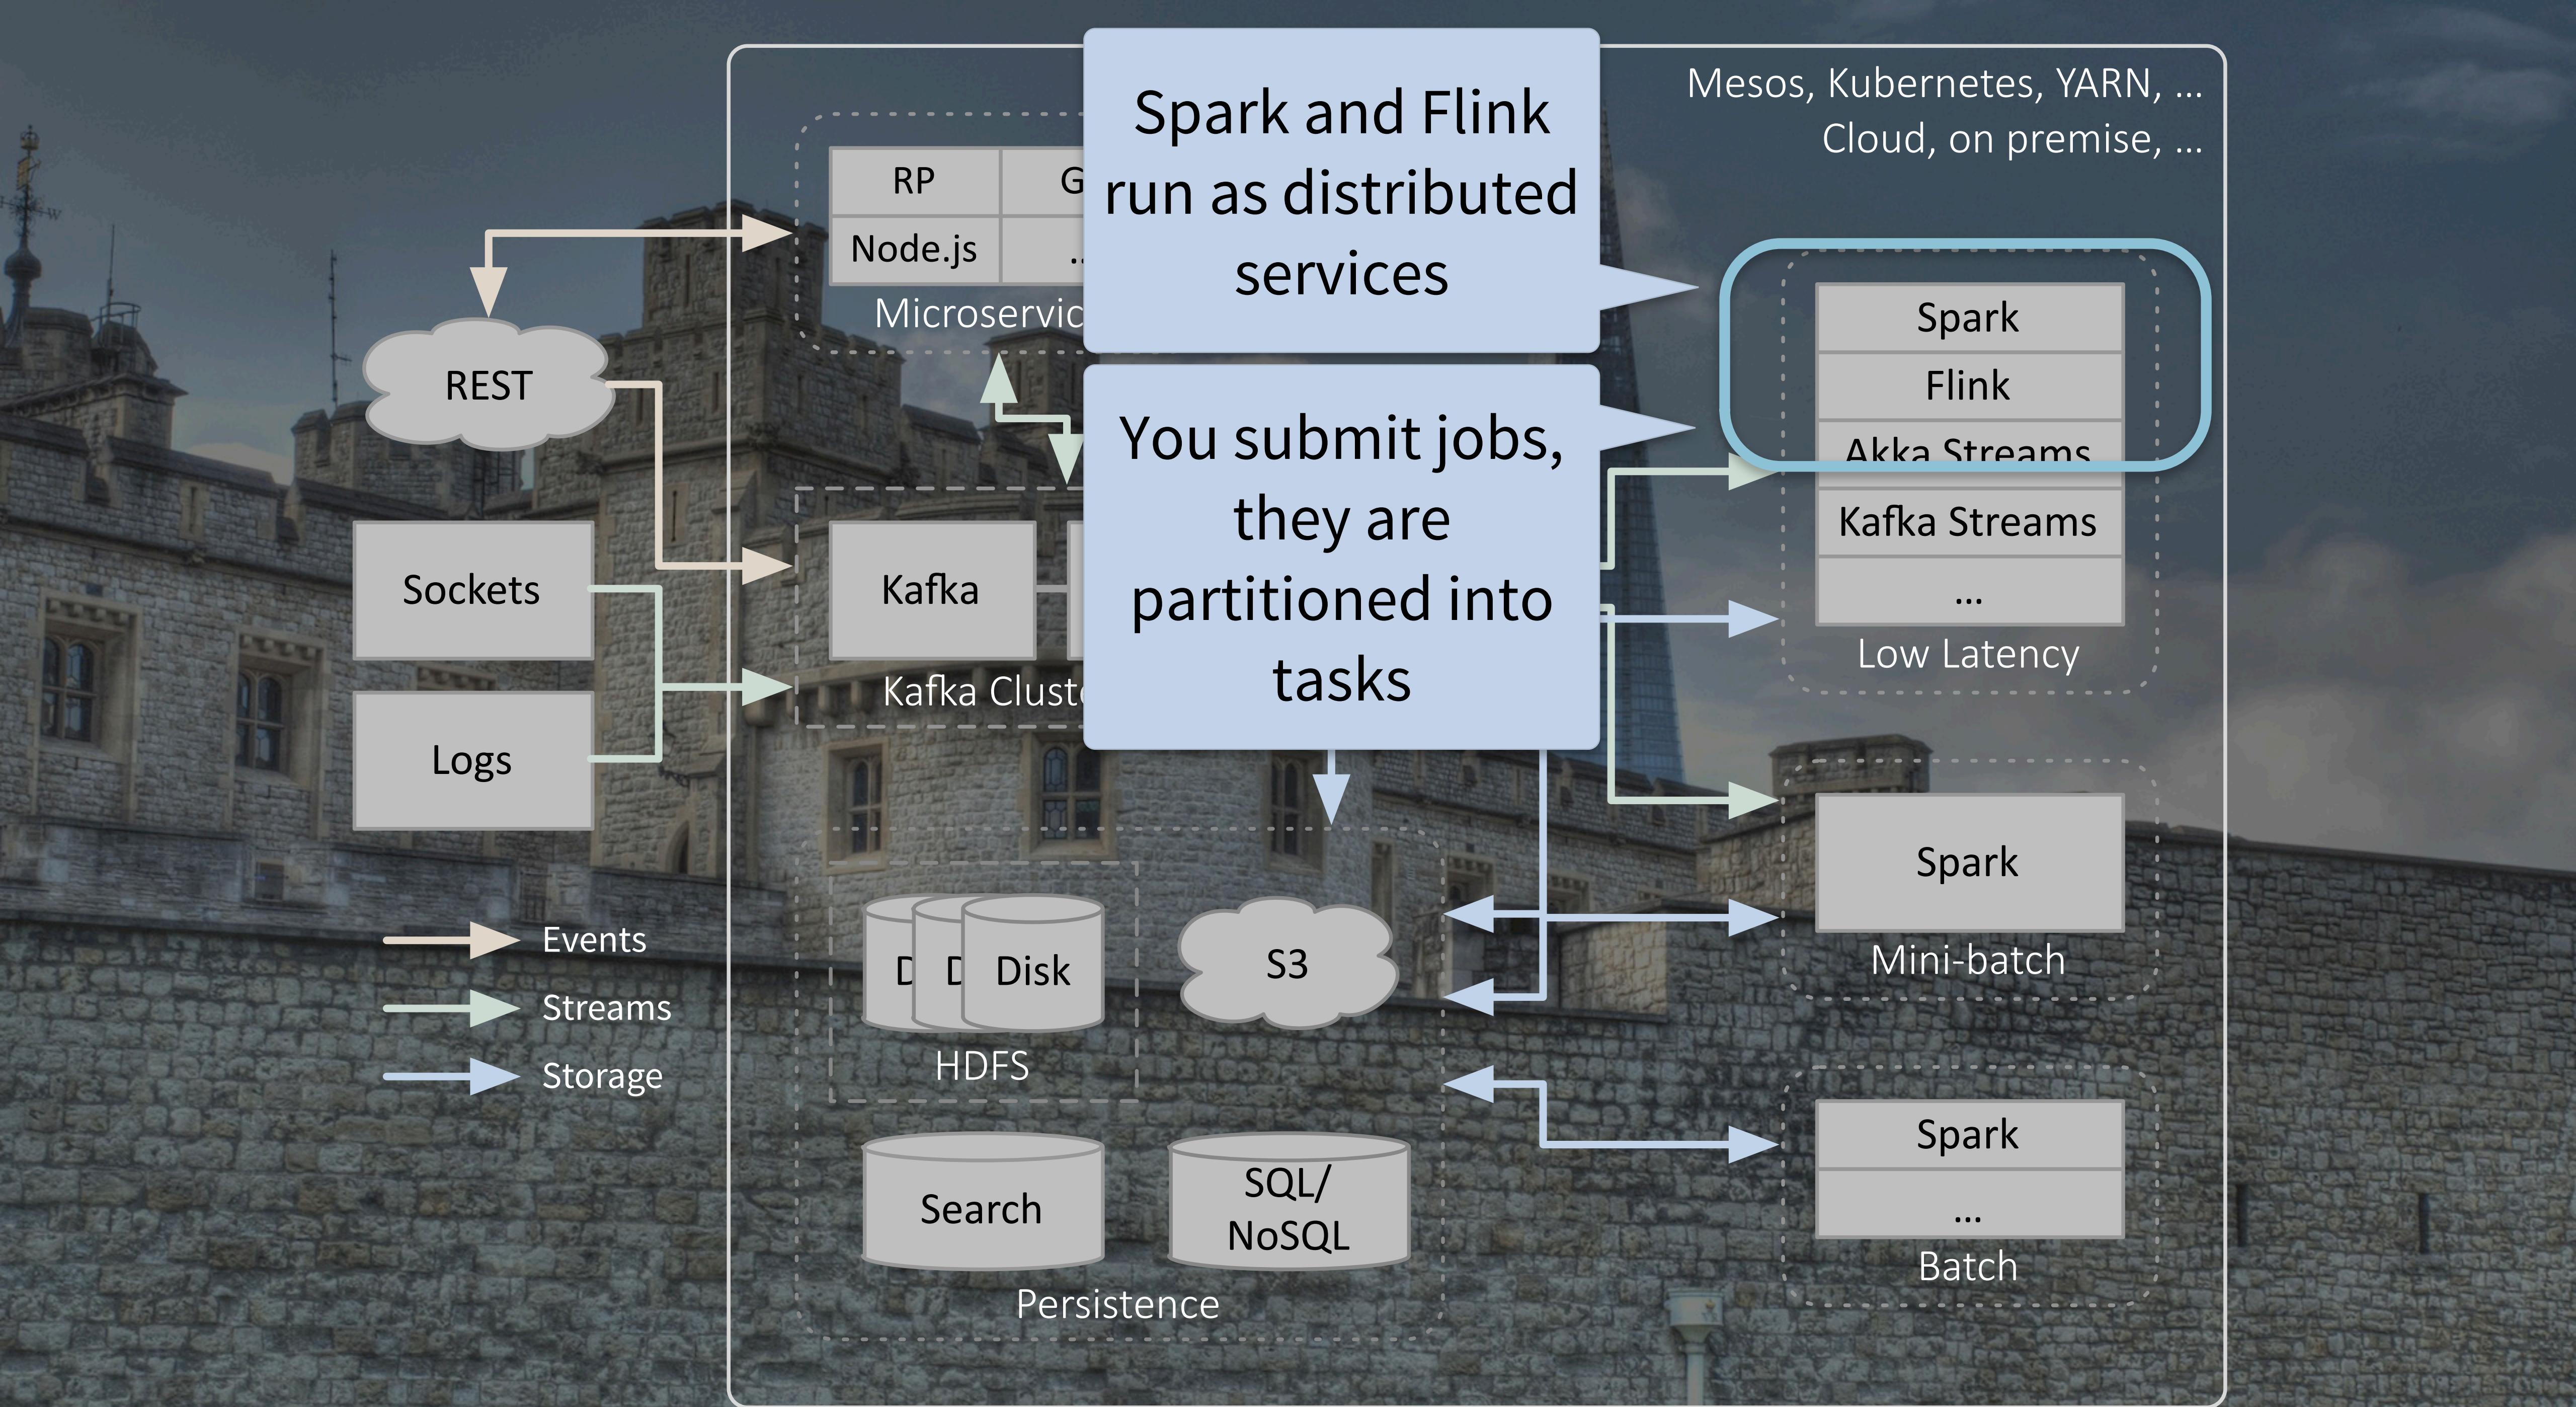
Kafka can simplify your architecture, too. Not only is this cleaner, it's more robust, as the point-to-point connections are more fragile, more likely to result in data and service loss if one point goes down, where Kafka keeps the other point "healthy" while the failed point is restored. Kafka also lets you easily support multiple consumers or producers per topic (the way data is organized, as in classic message queues). Finally, the uniformity of always (or most of the time) communicating through Kafka simplifies the challenge of connecting services together with different APIs.



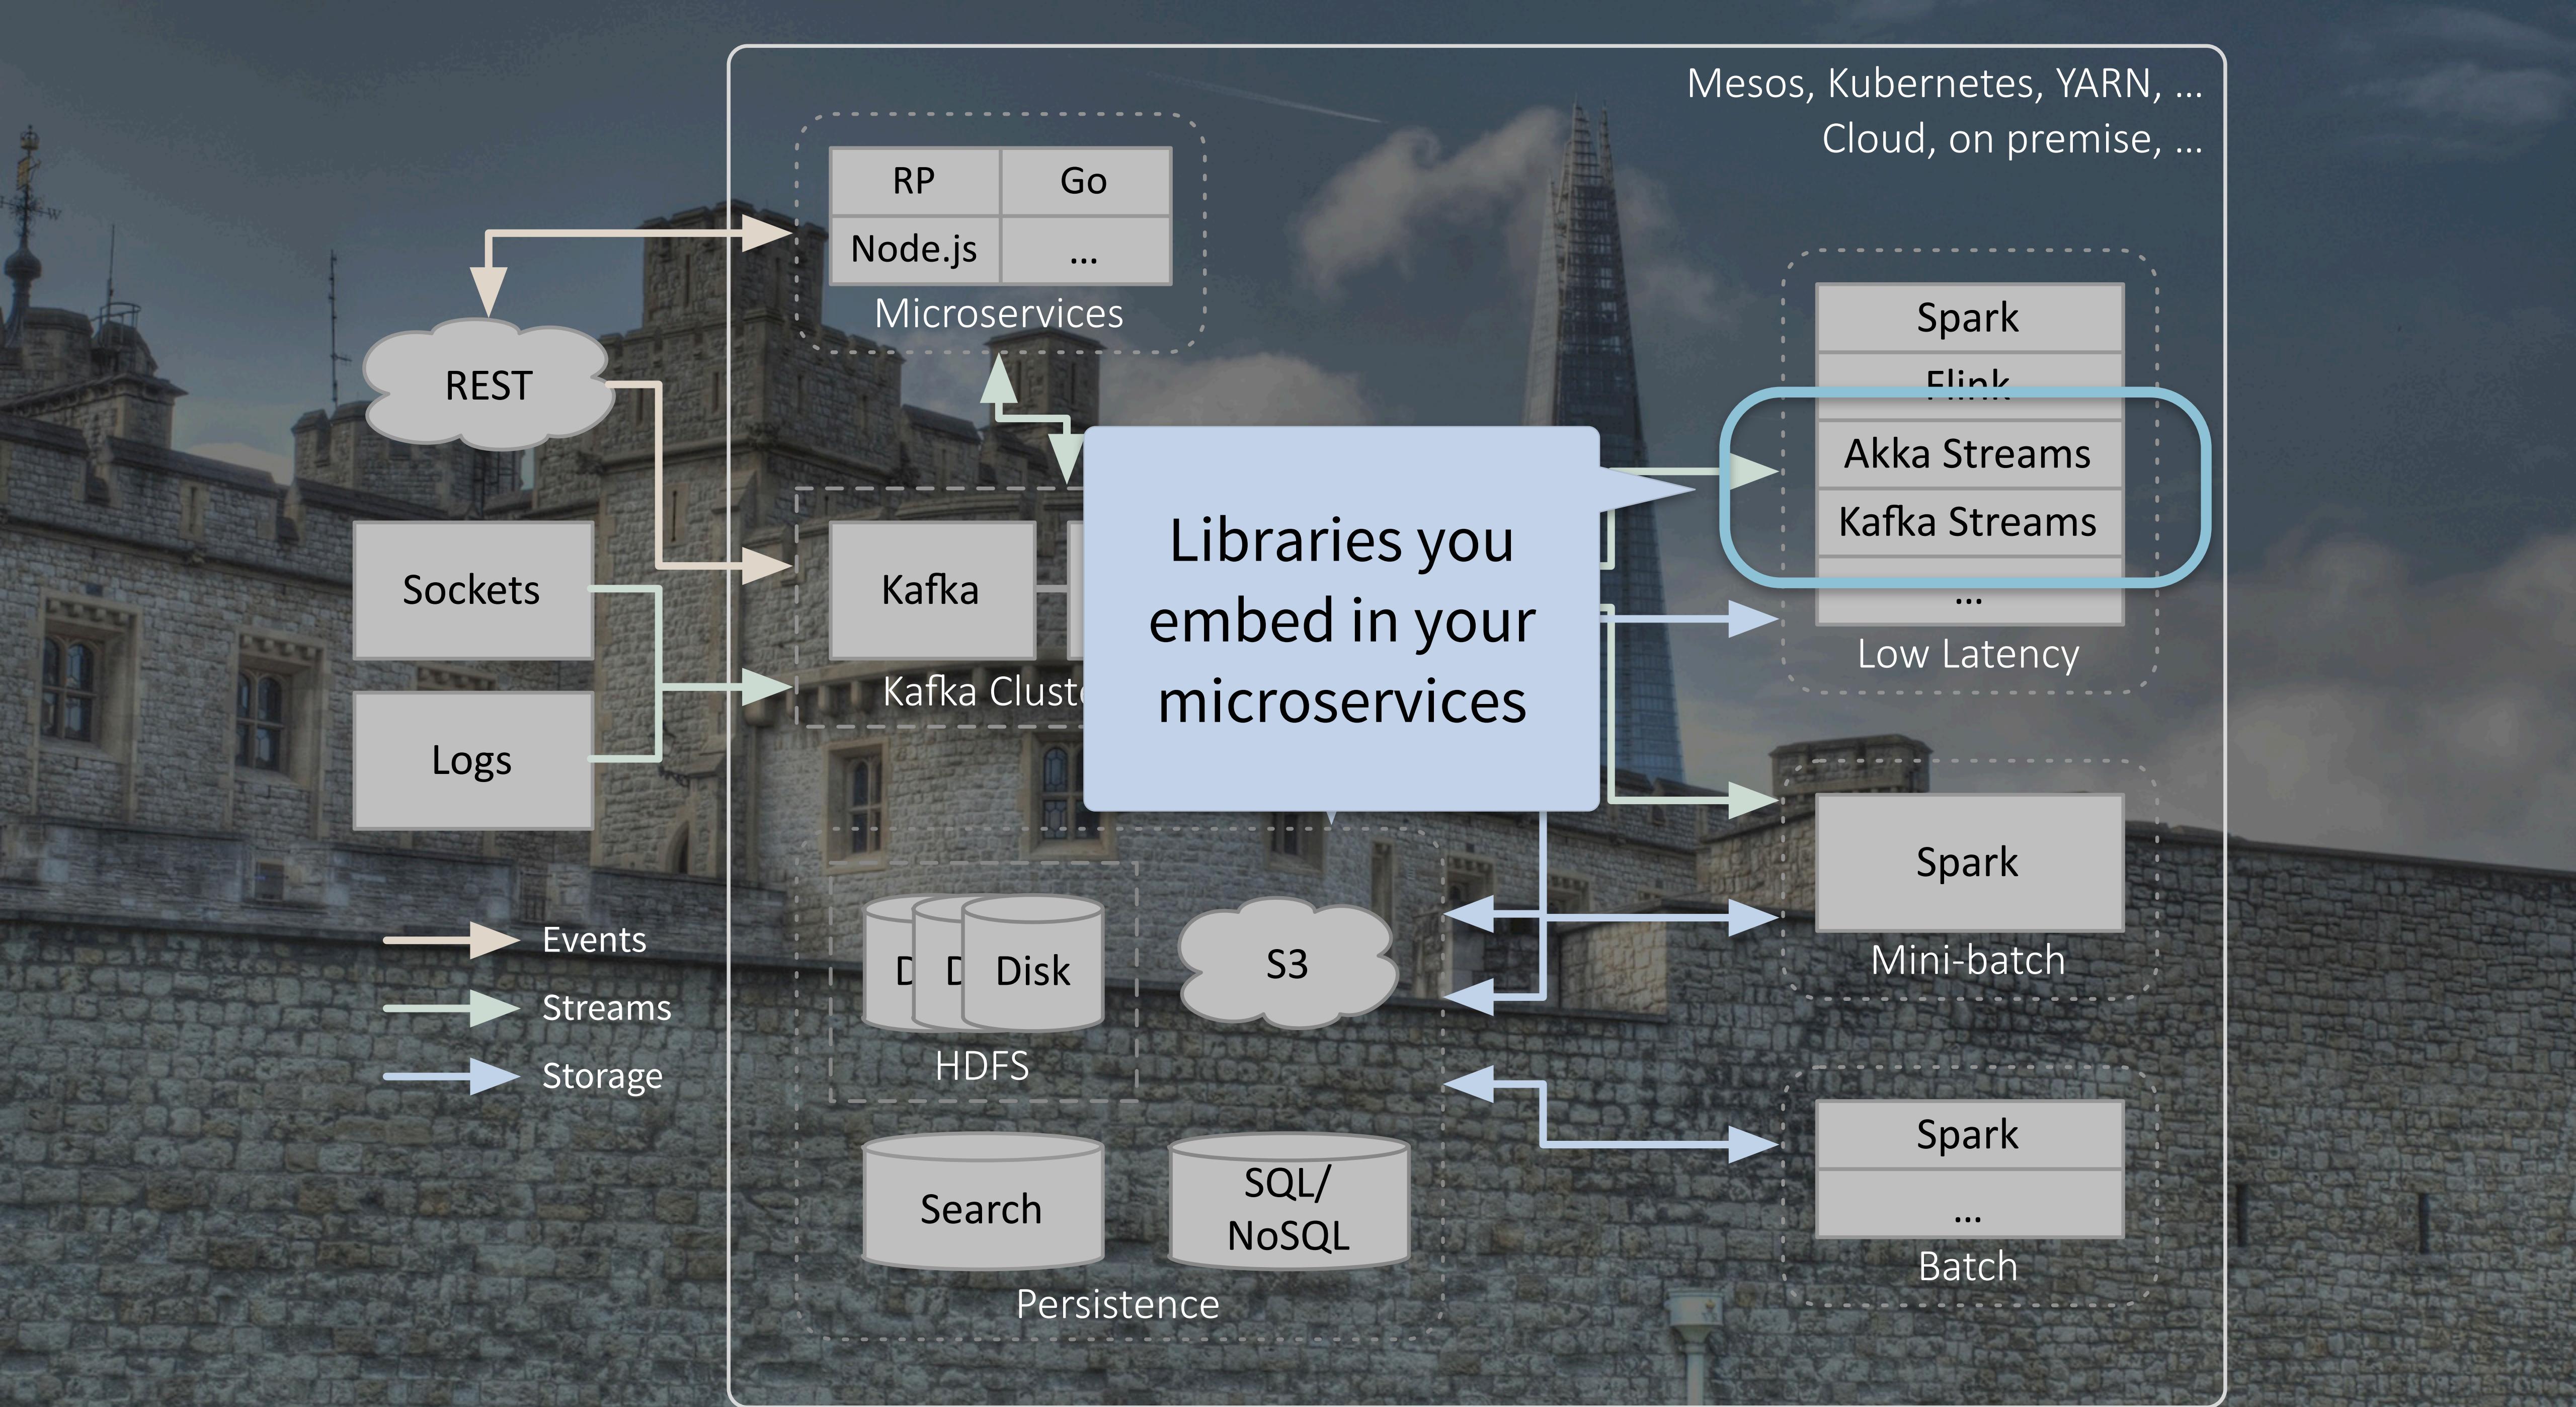
How do you choose?

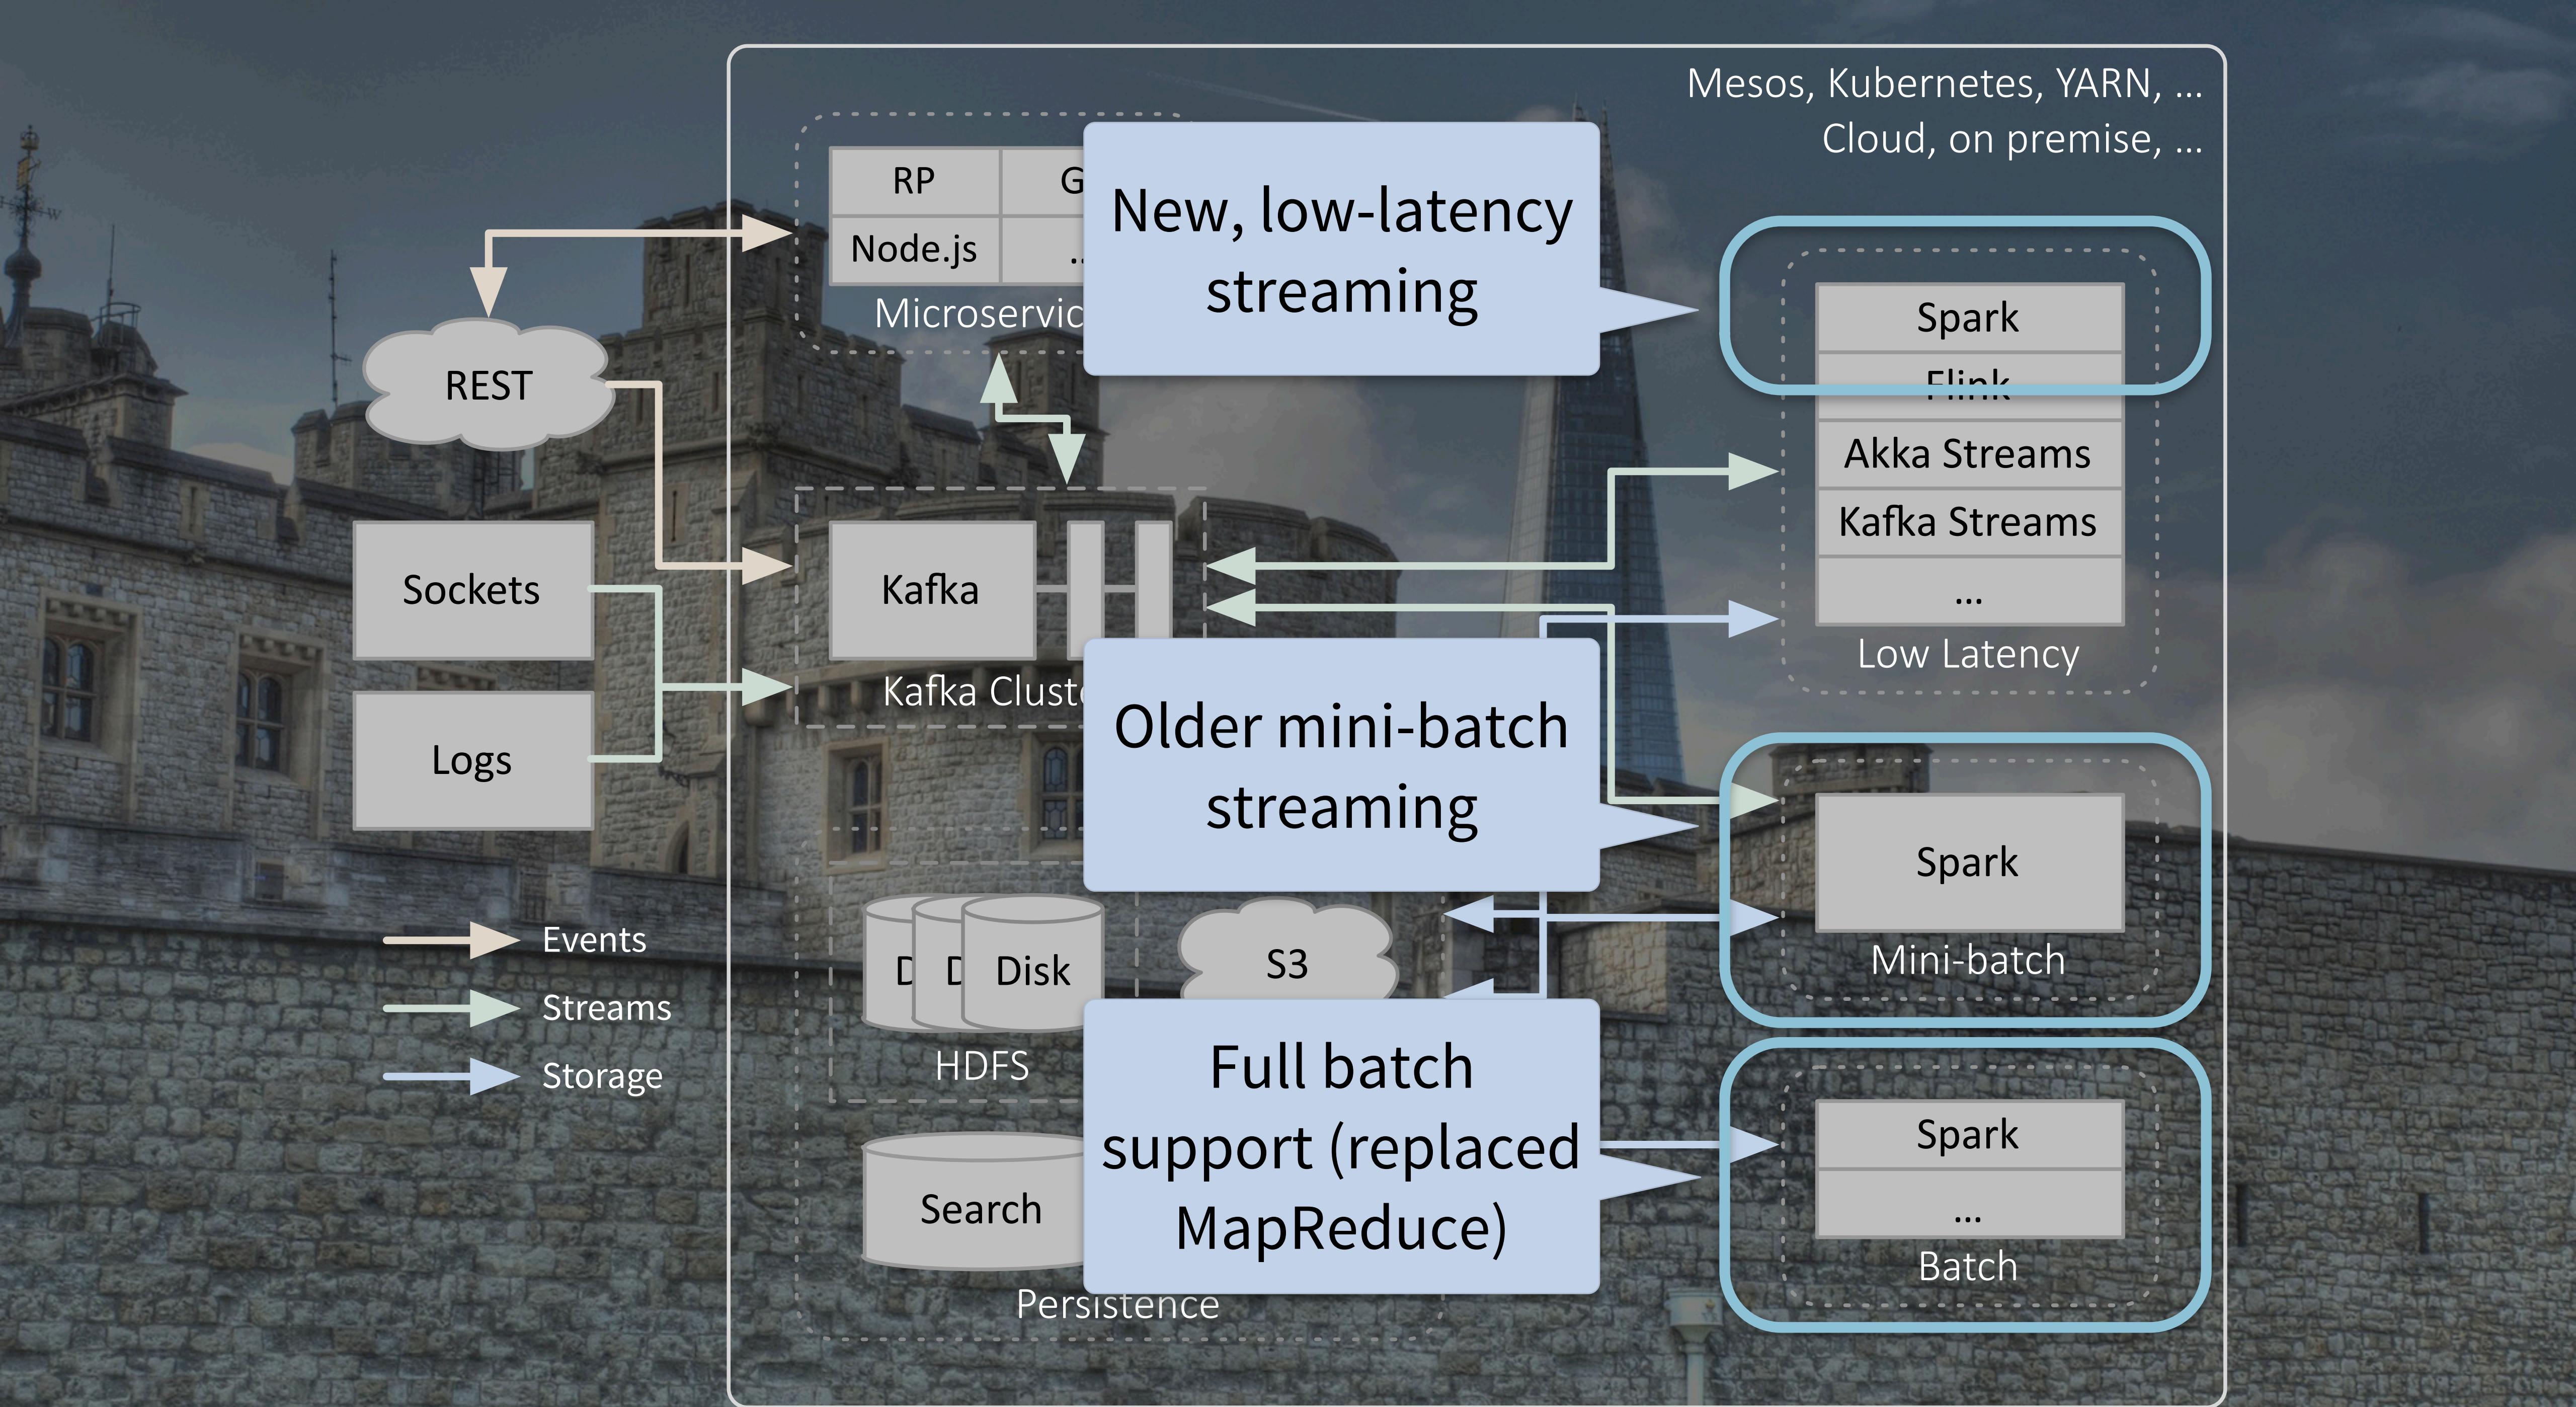
- Latency? How low?
- Volume: How high?
- Which kinds of data processing?
- What's your preferred application architecture?

If you have tight latency requirements, you can't use a mini-batch or batch engine. If your latency constraints are more flexible, you have more tool options and you can do more sophisticated and expensive things, like training ML models, write to databases, etc.
What's your volume? Do you need to partition your streams over a cluster to use parallelism to handle the load?
Are you running SQL-like queries over the data? Are you doing "simple" ETL? Are you training ML models or serving models?
Do you want to run microservices for data, too? Is the overhead and compute "model" of Spark or Flink right for you, or do you need more flexibility?

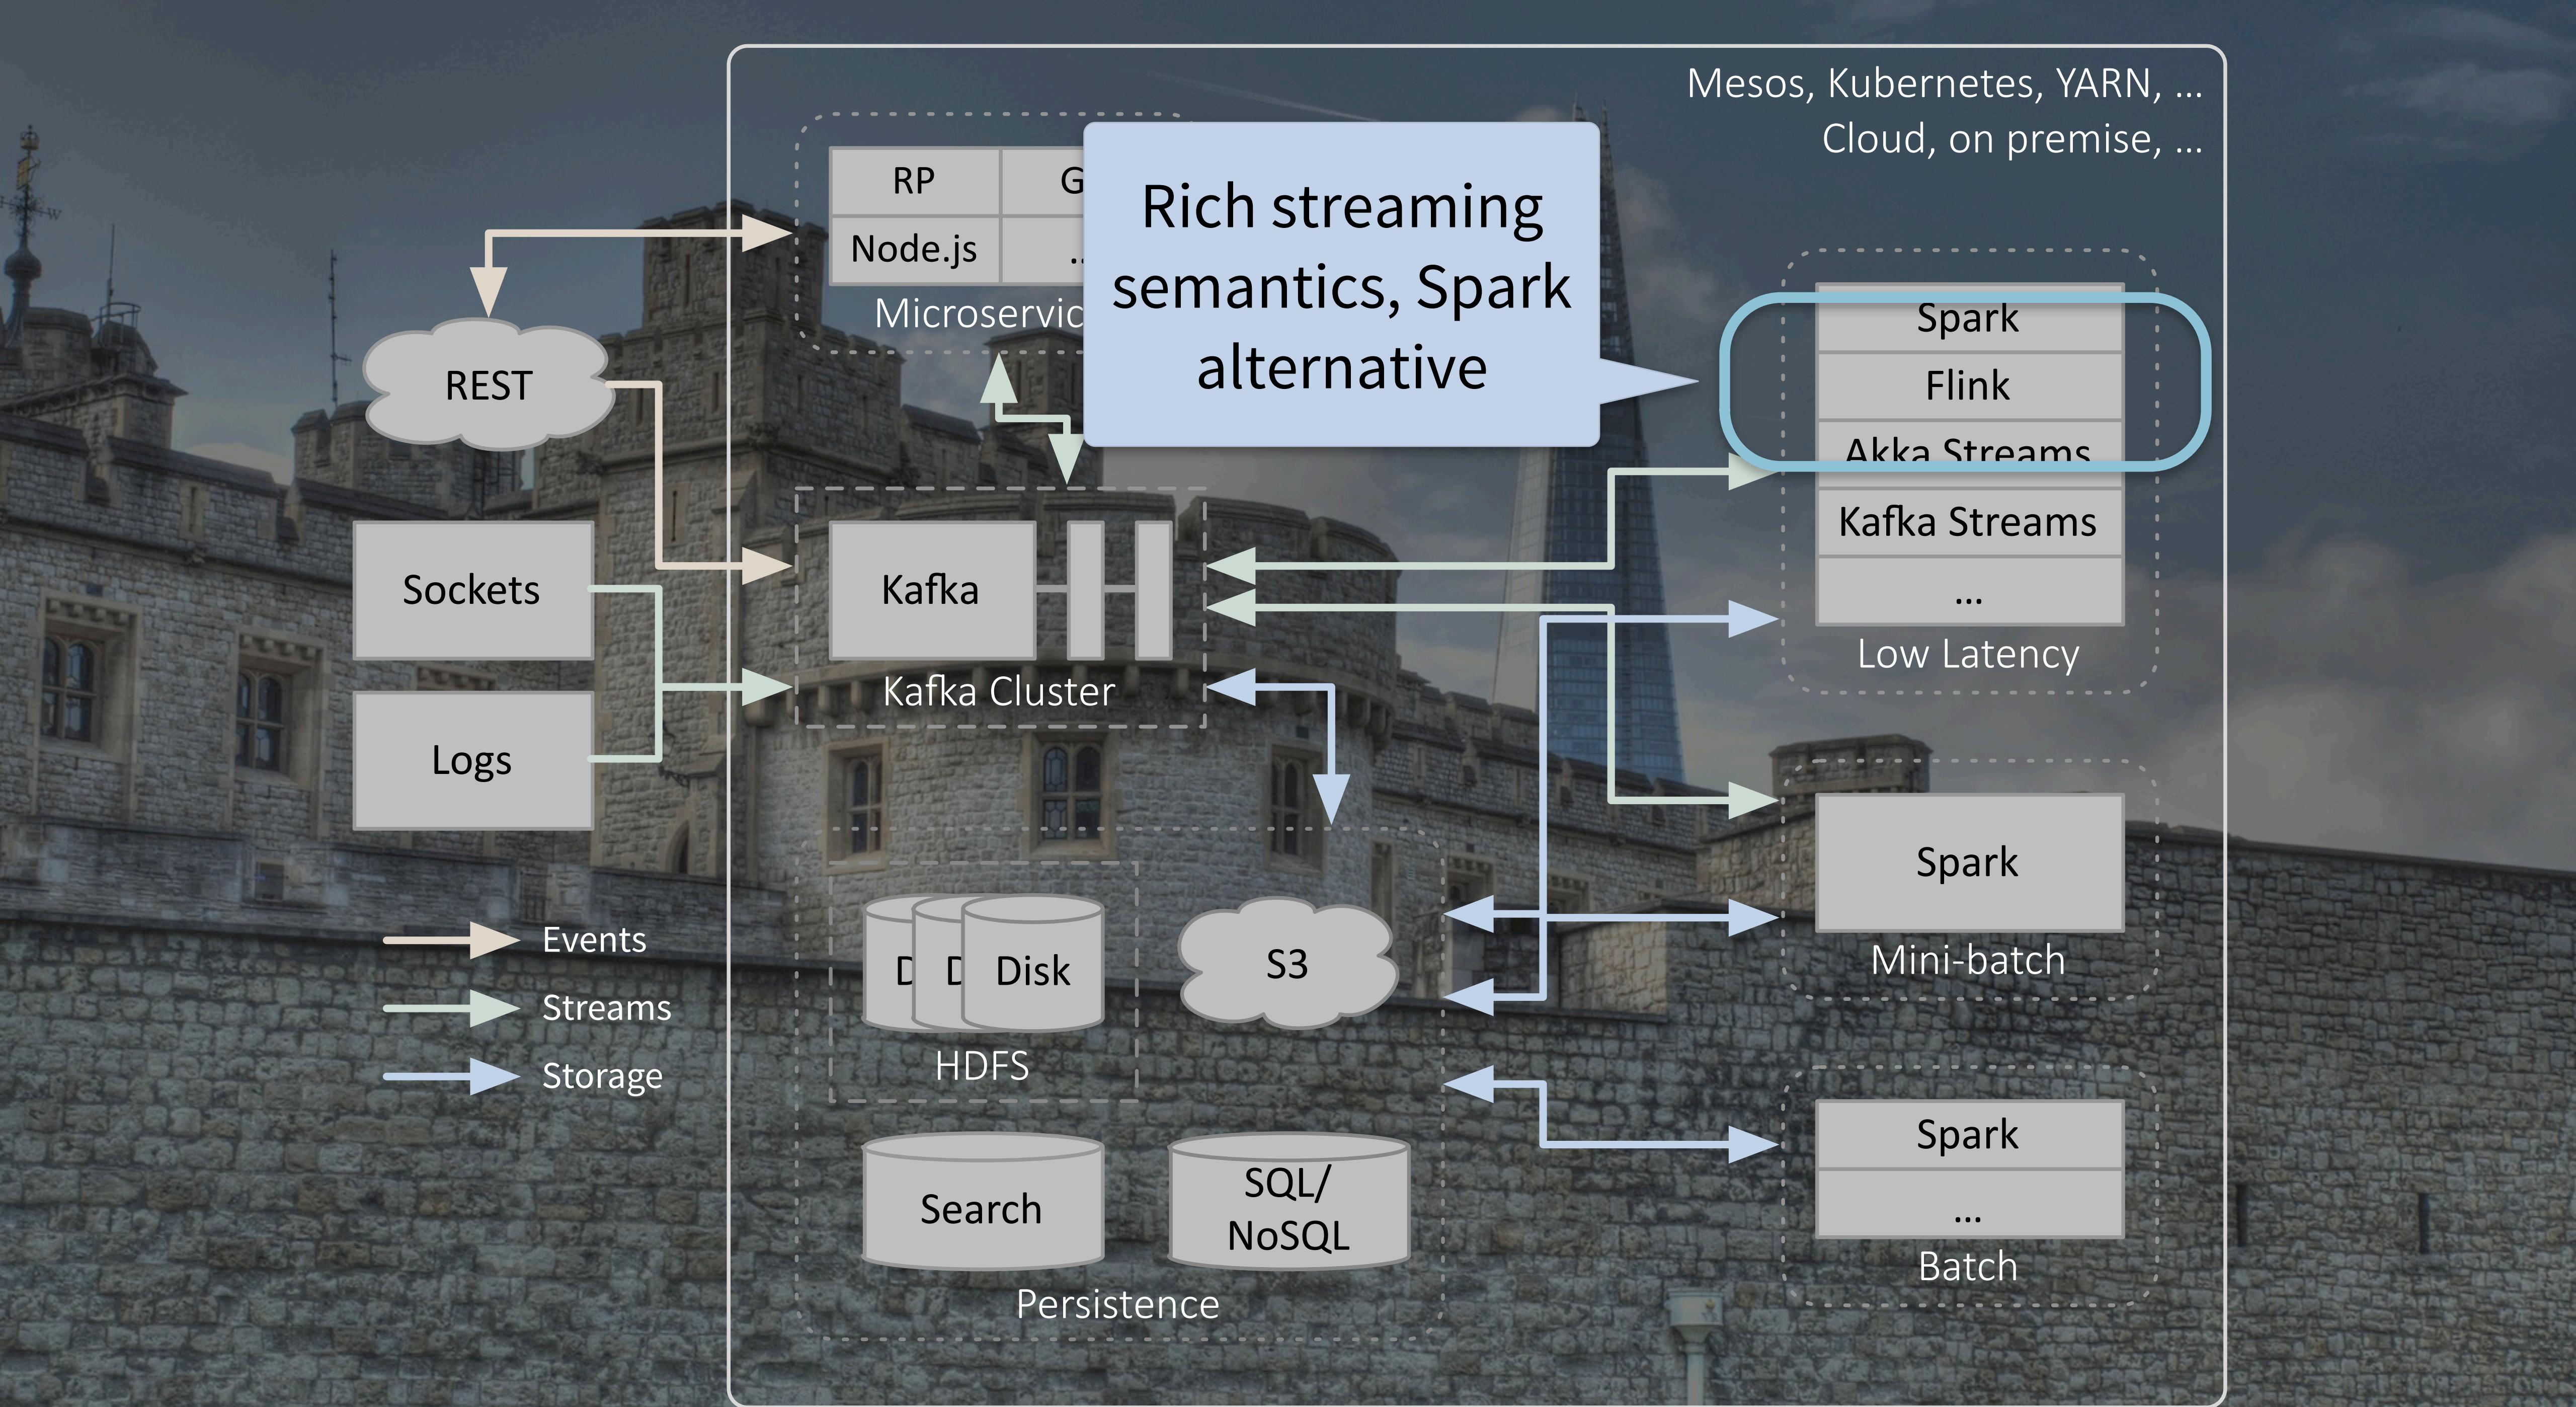


The best known is Spark, which provides rich SQL and dataflow APIs that support batch jobs and streaming jobs. There are two streaming APIs, an older “mini-batch” model, where data is captured in time intervals down to ~100 milliseconds in size, then processed as a small batch. Now there is a new low-latency API.

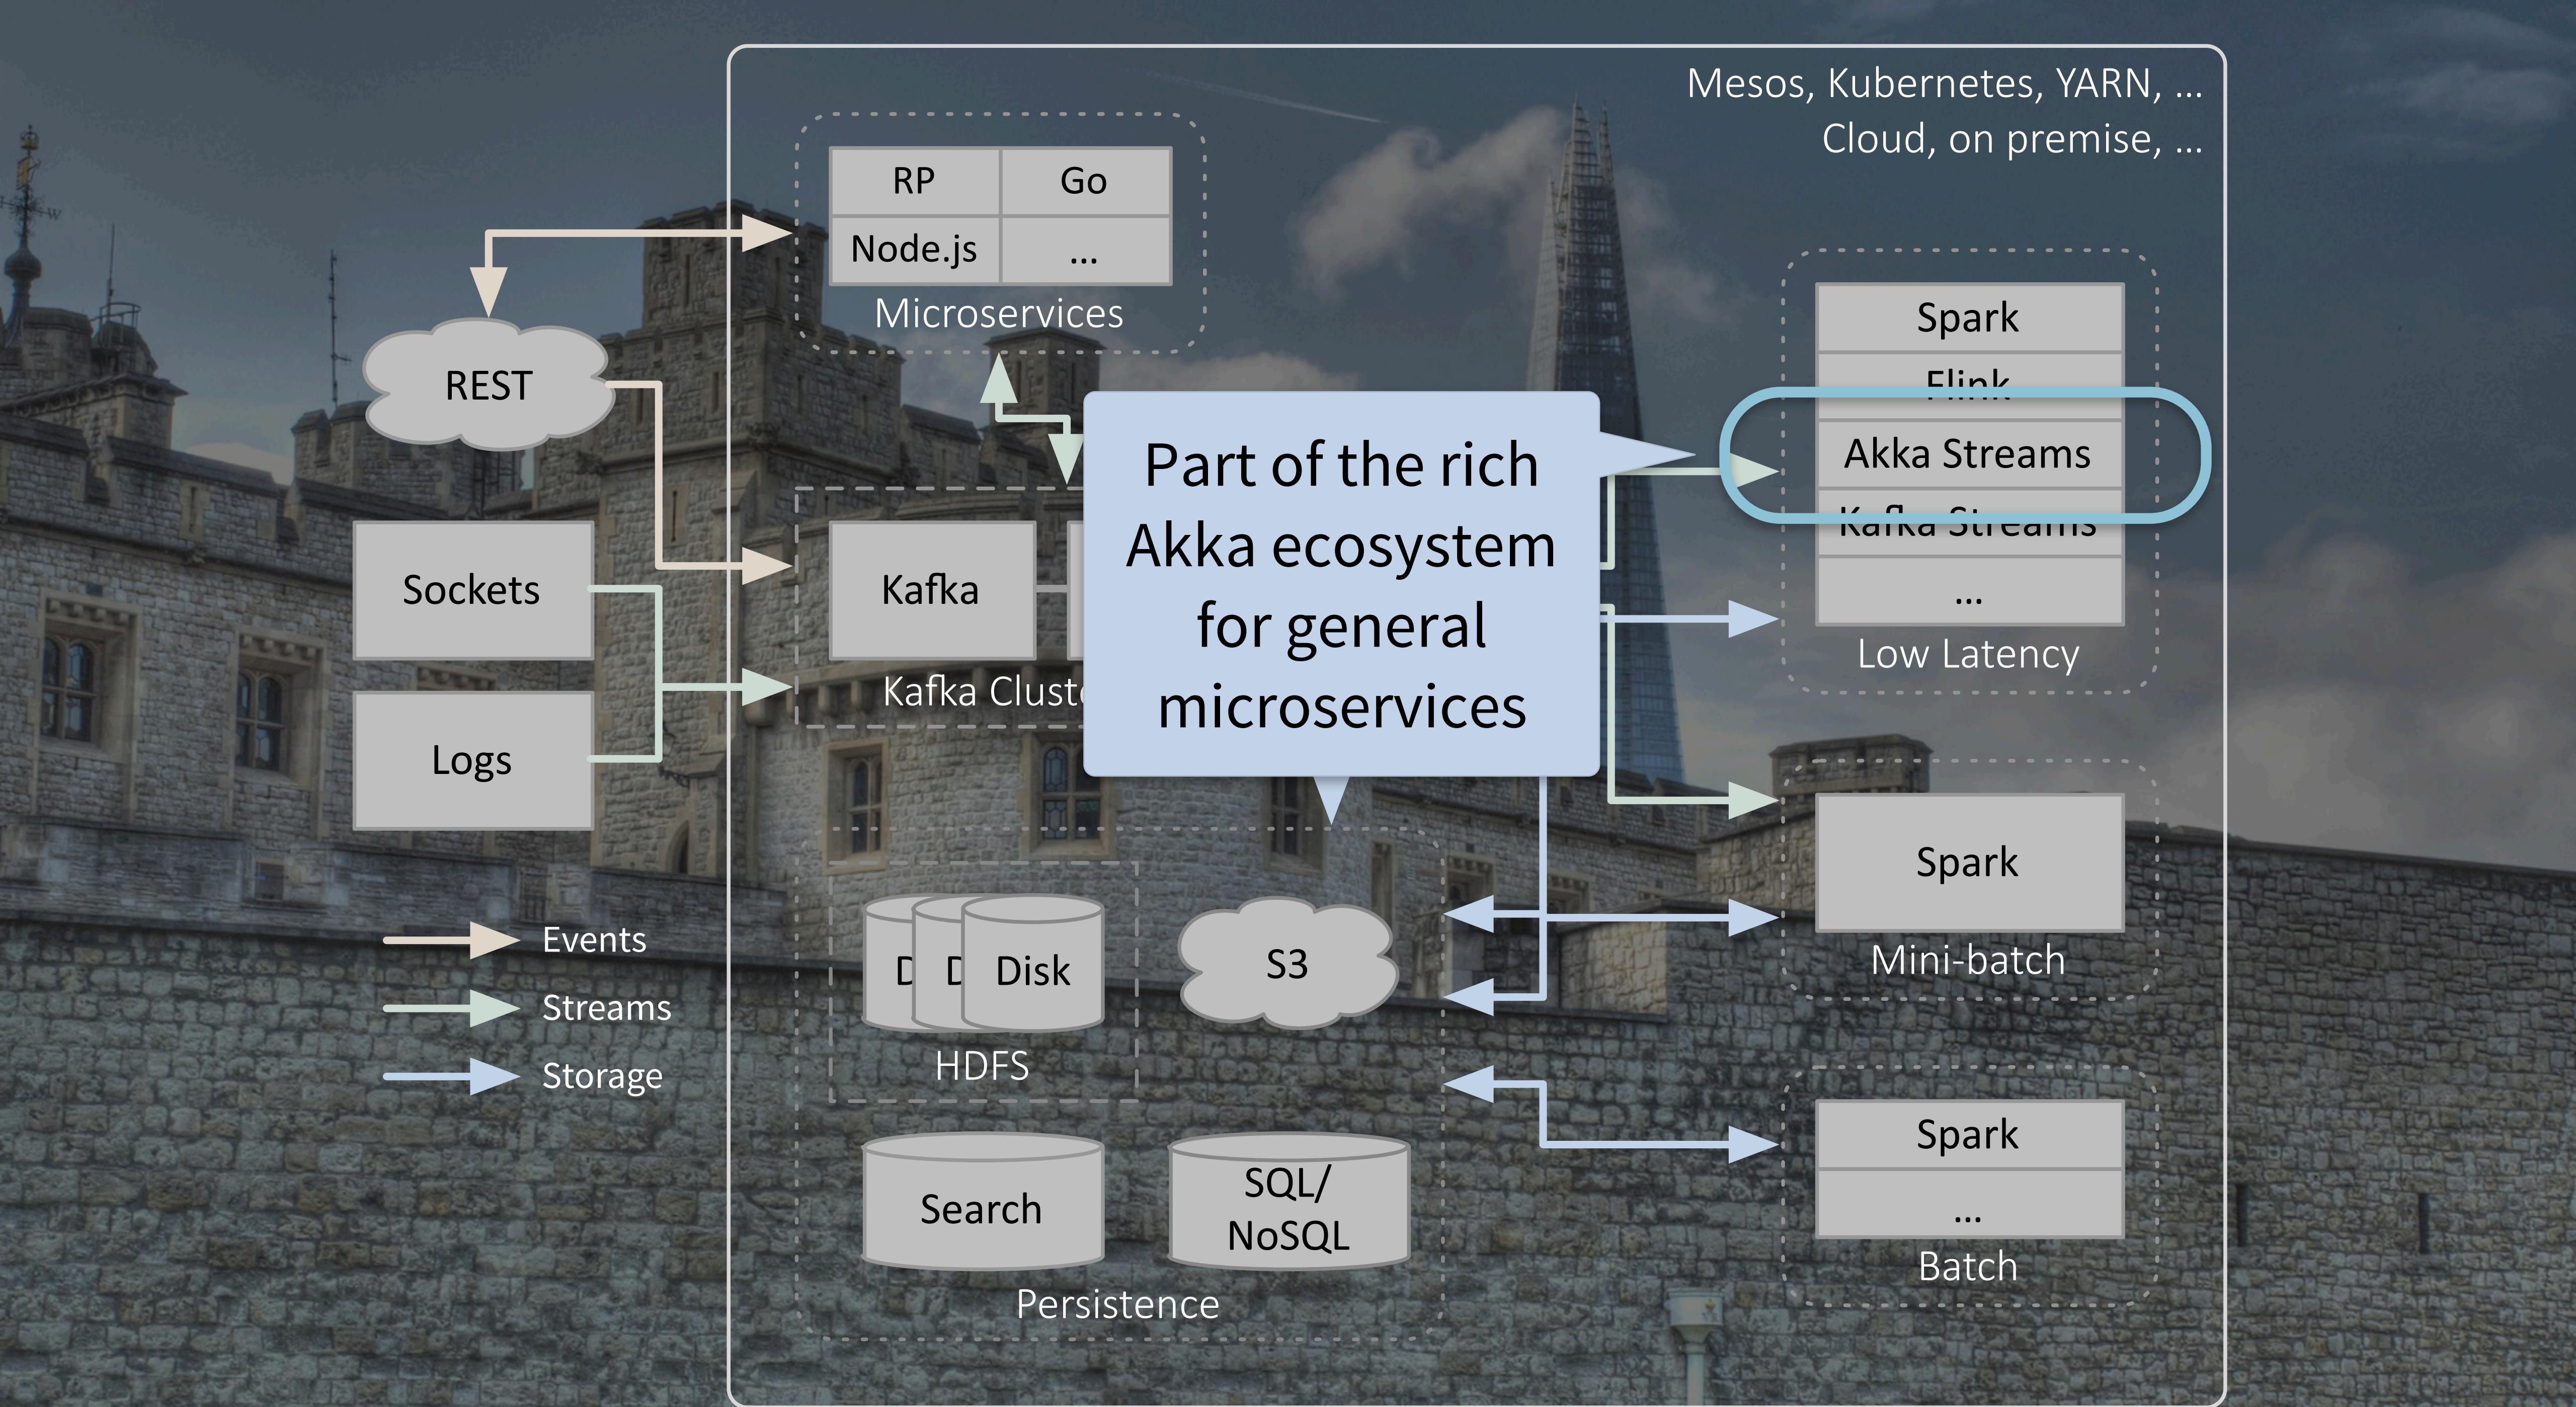


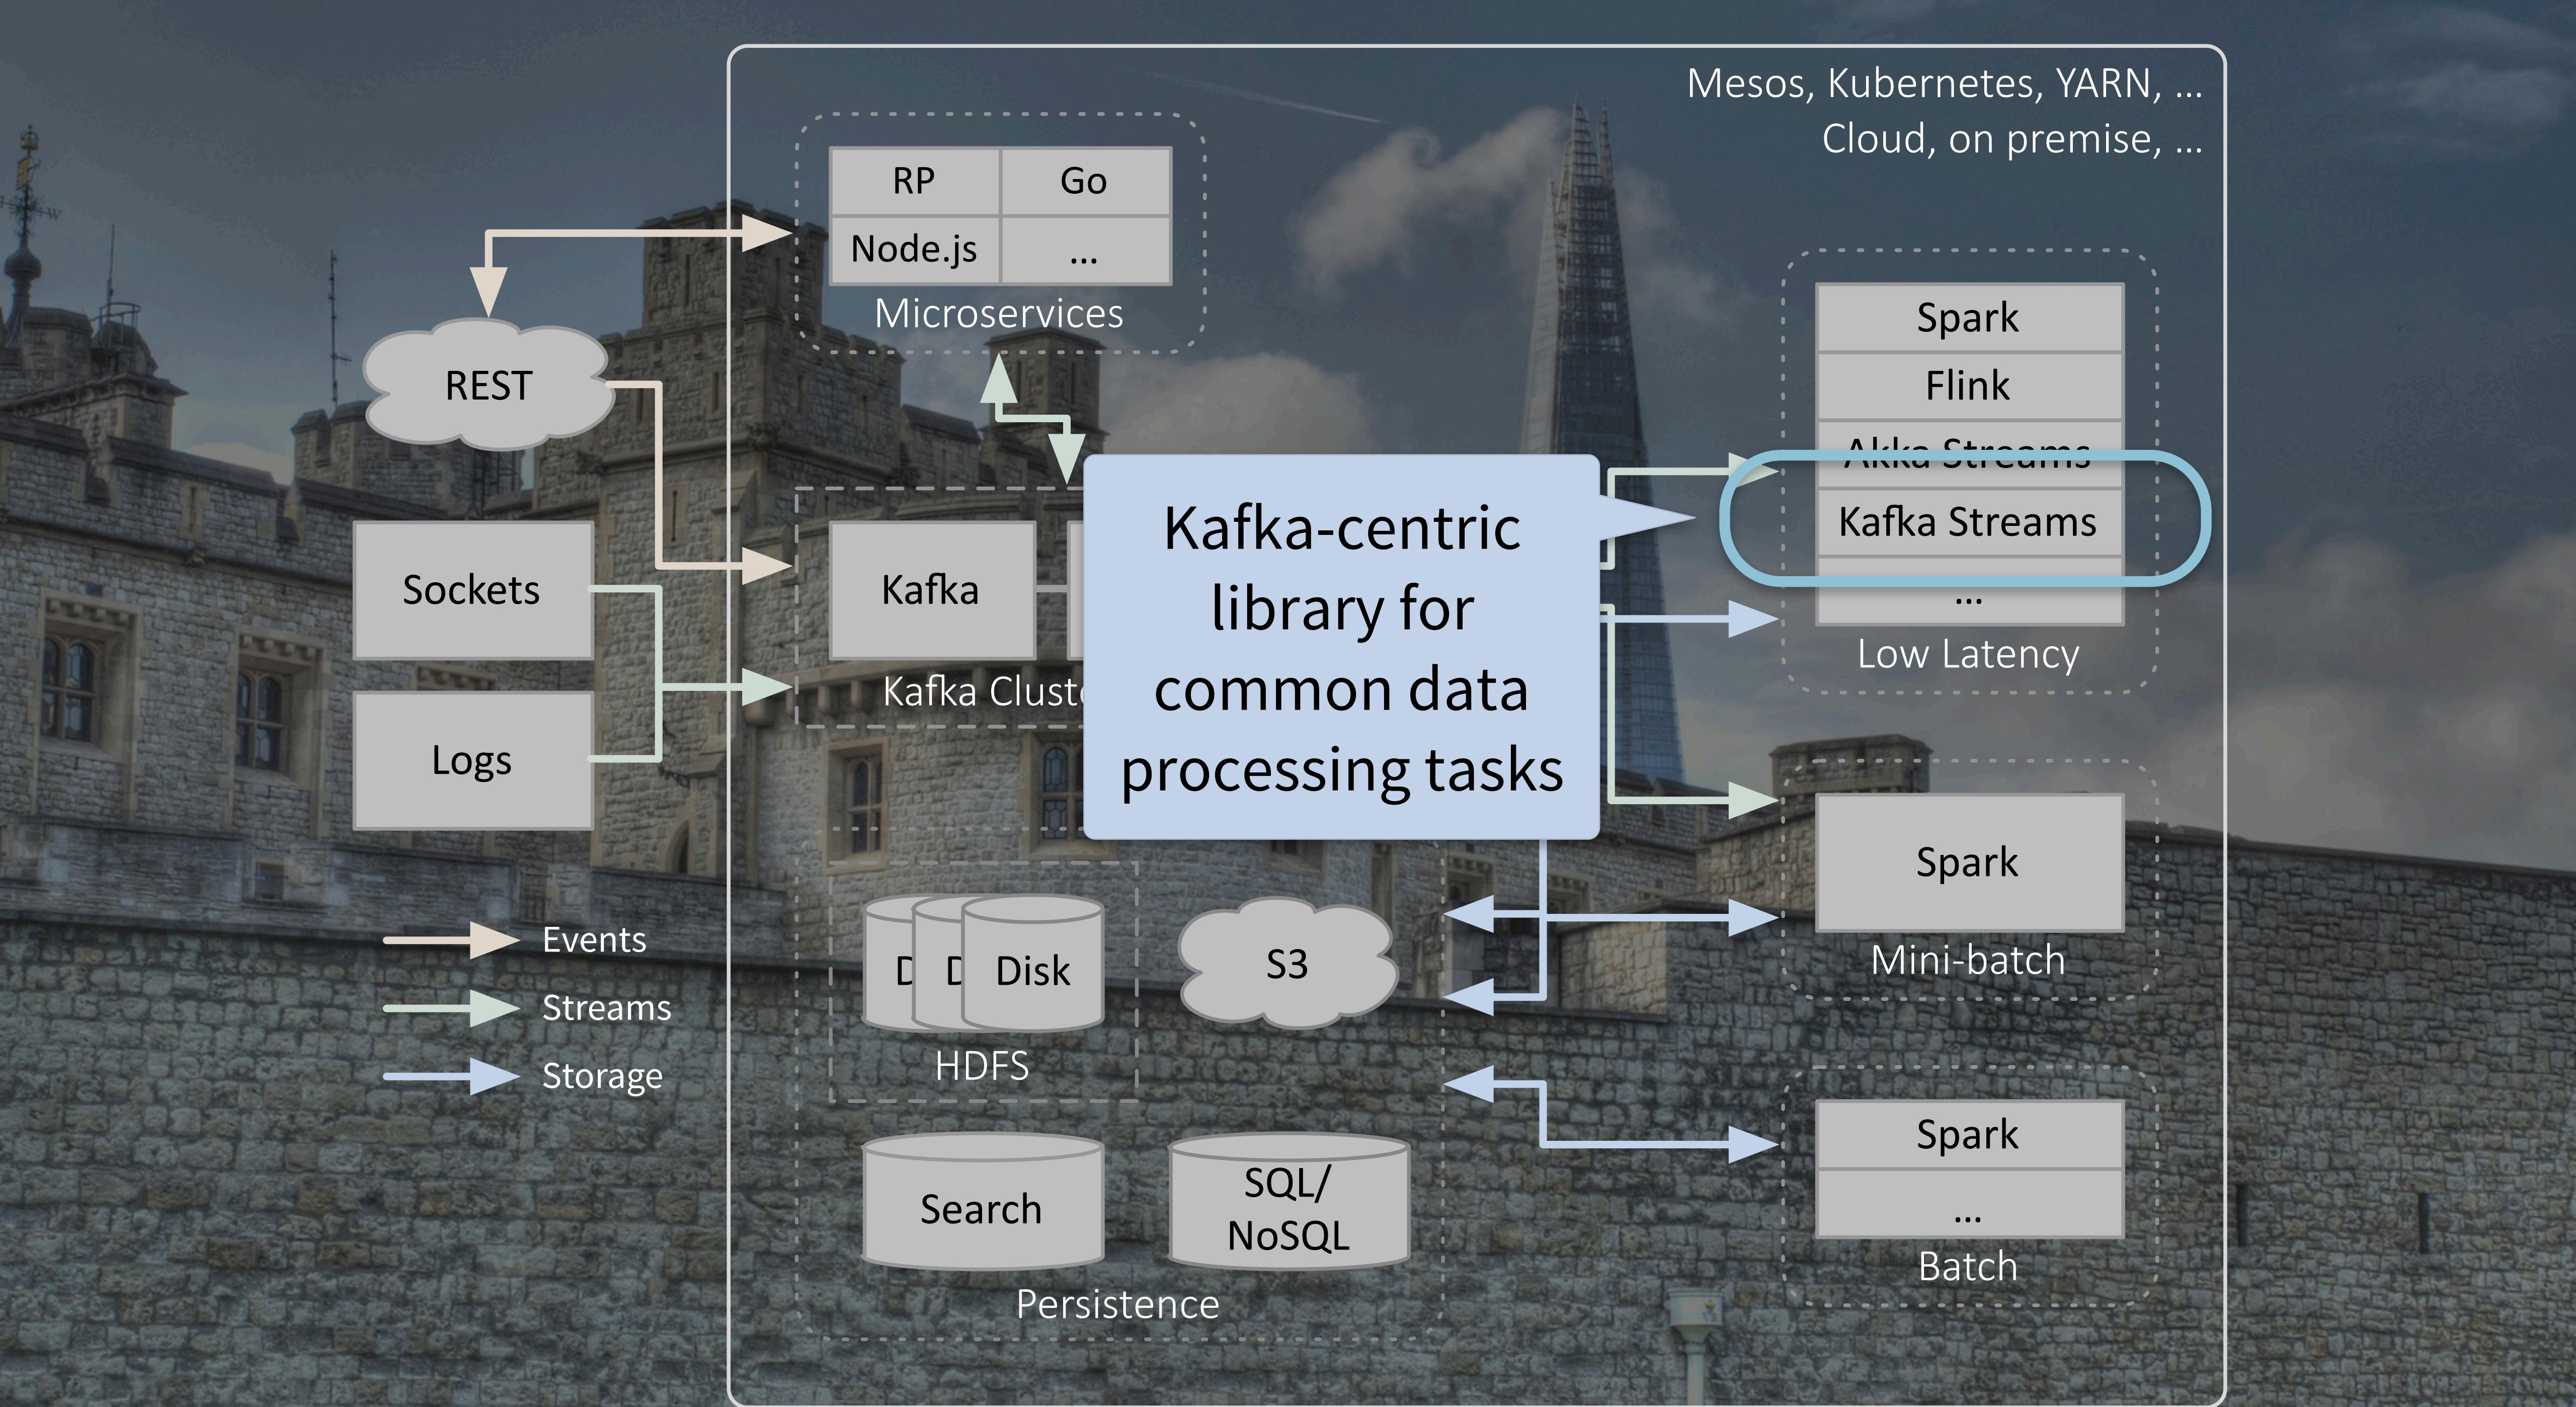


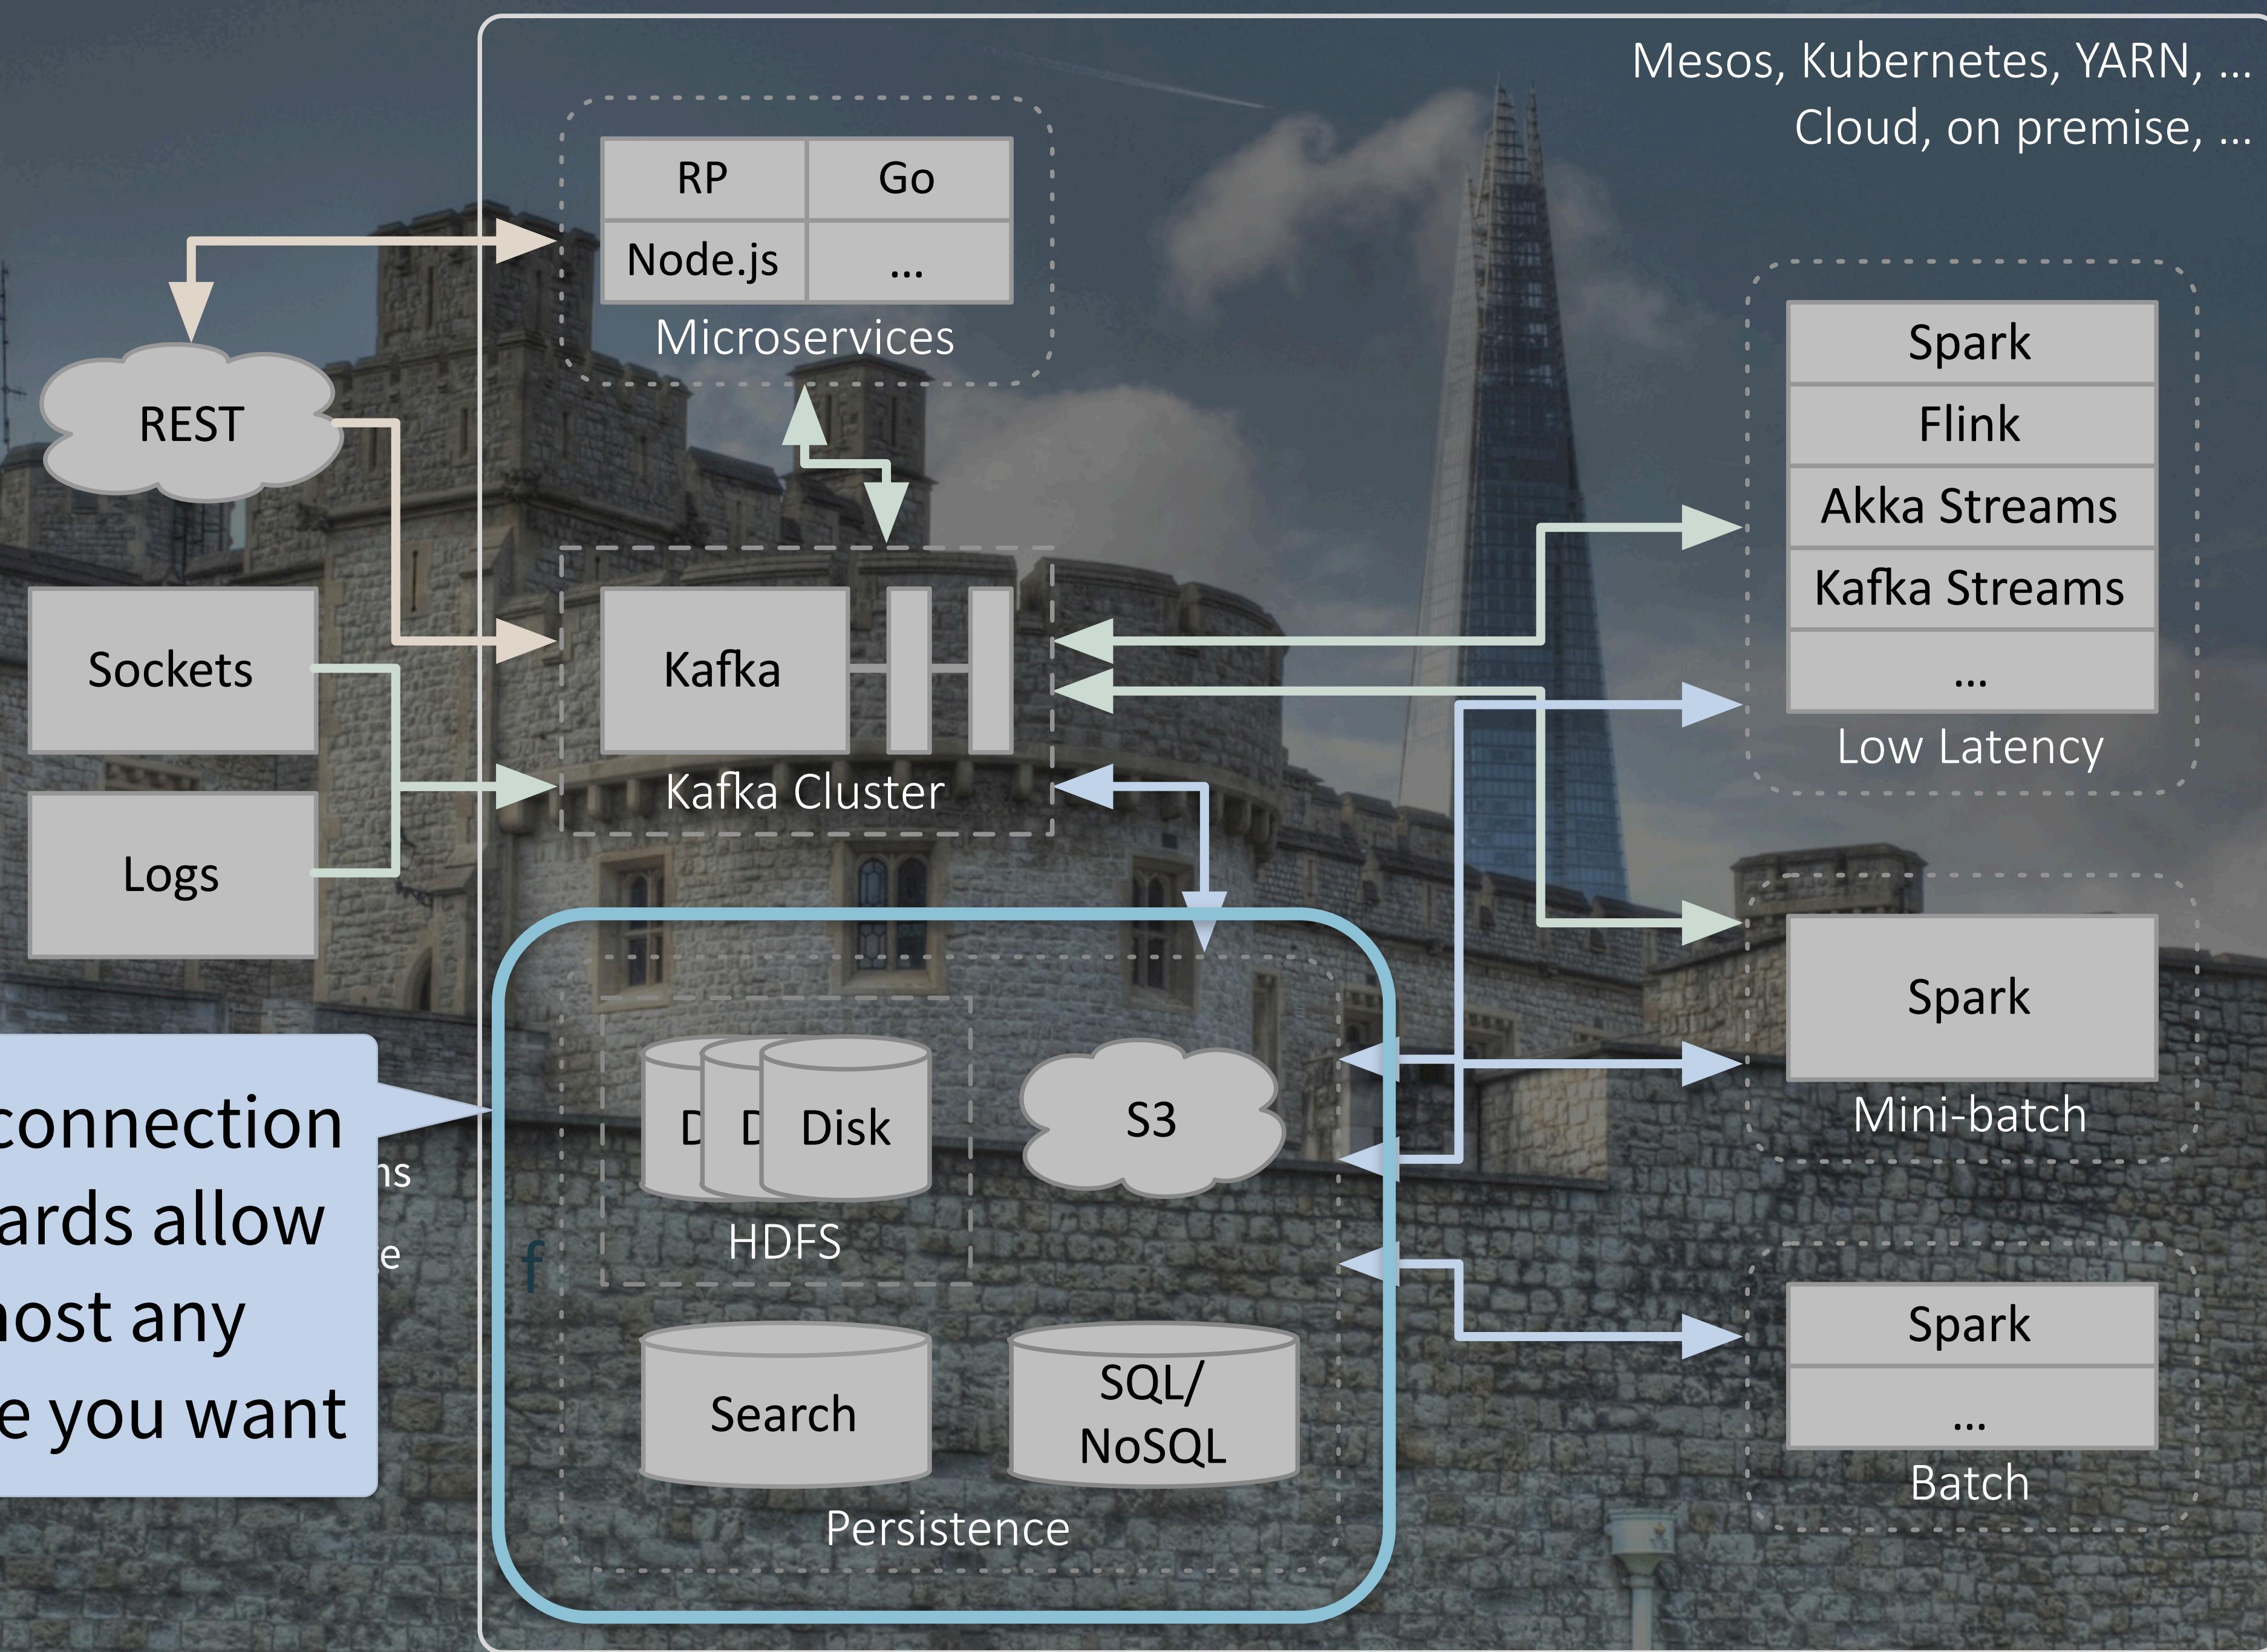
The best known is Spark, which provides rich SQL and dataflow APIs that support batch jobs and streaming jobs. There are two streaming APIs, an older “mini-batch” model, where data is captured in time intervals down to ~100 milliseconds in size, then processed as a small batch. Now there is a new low-latency API.



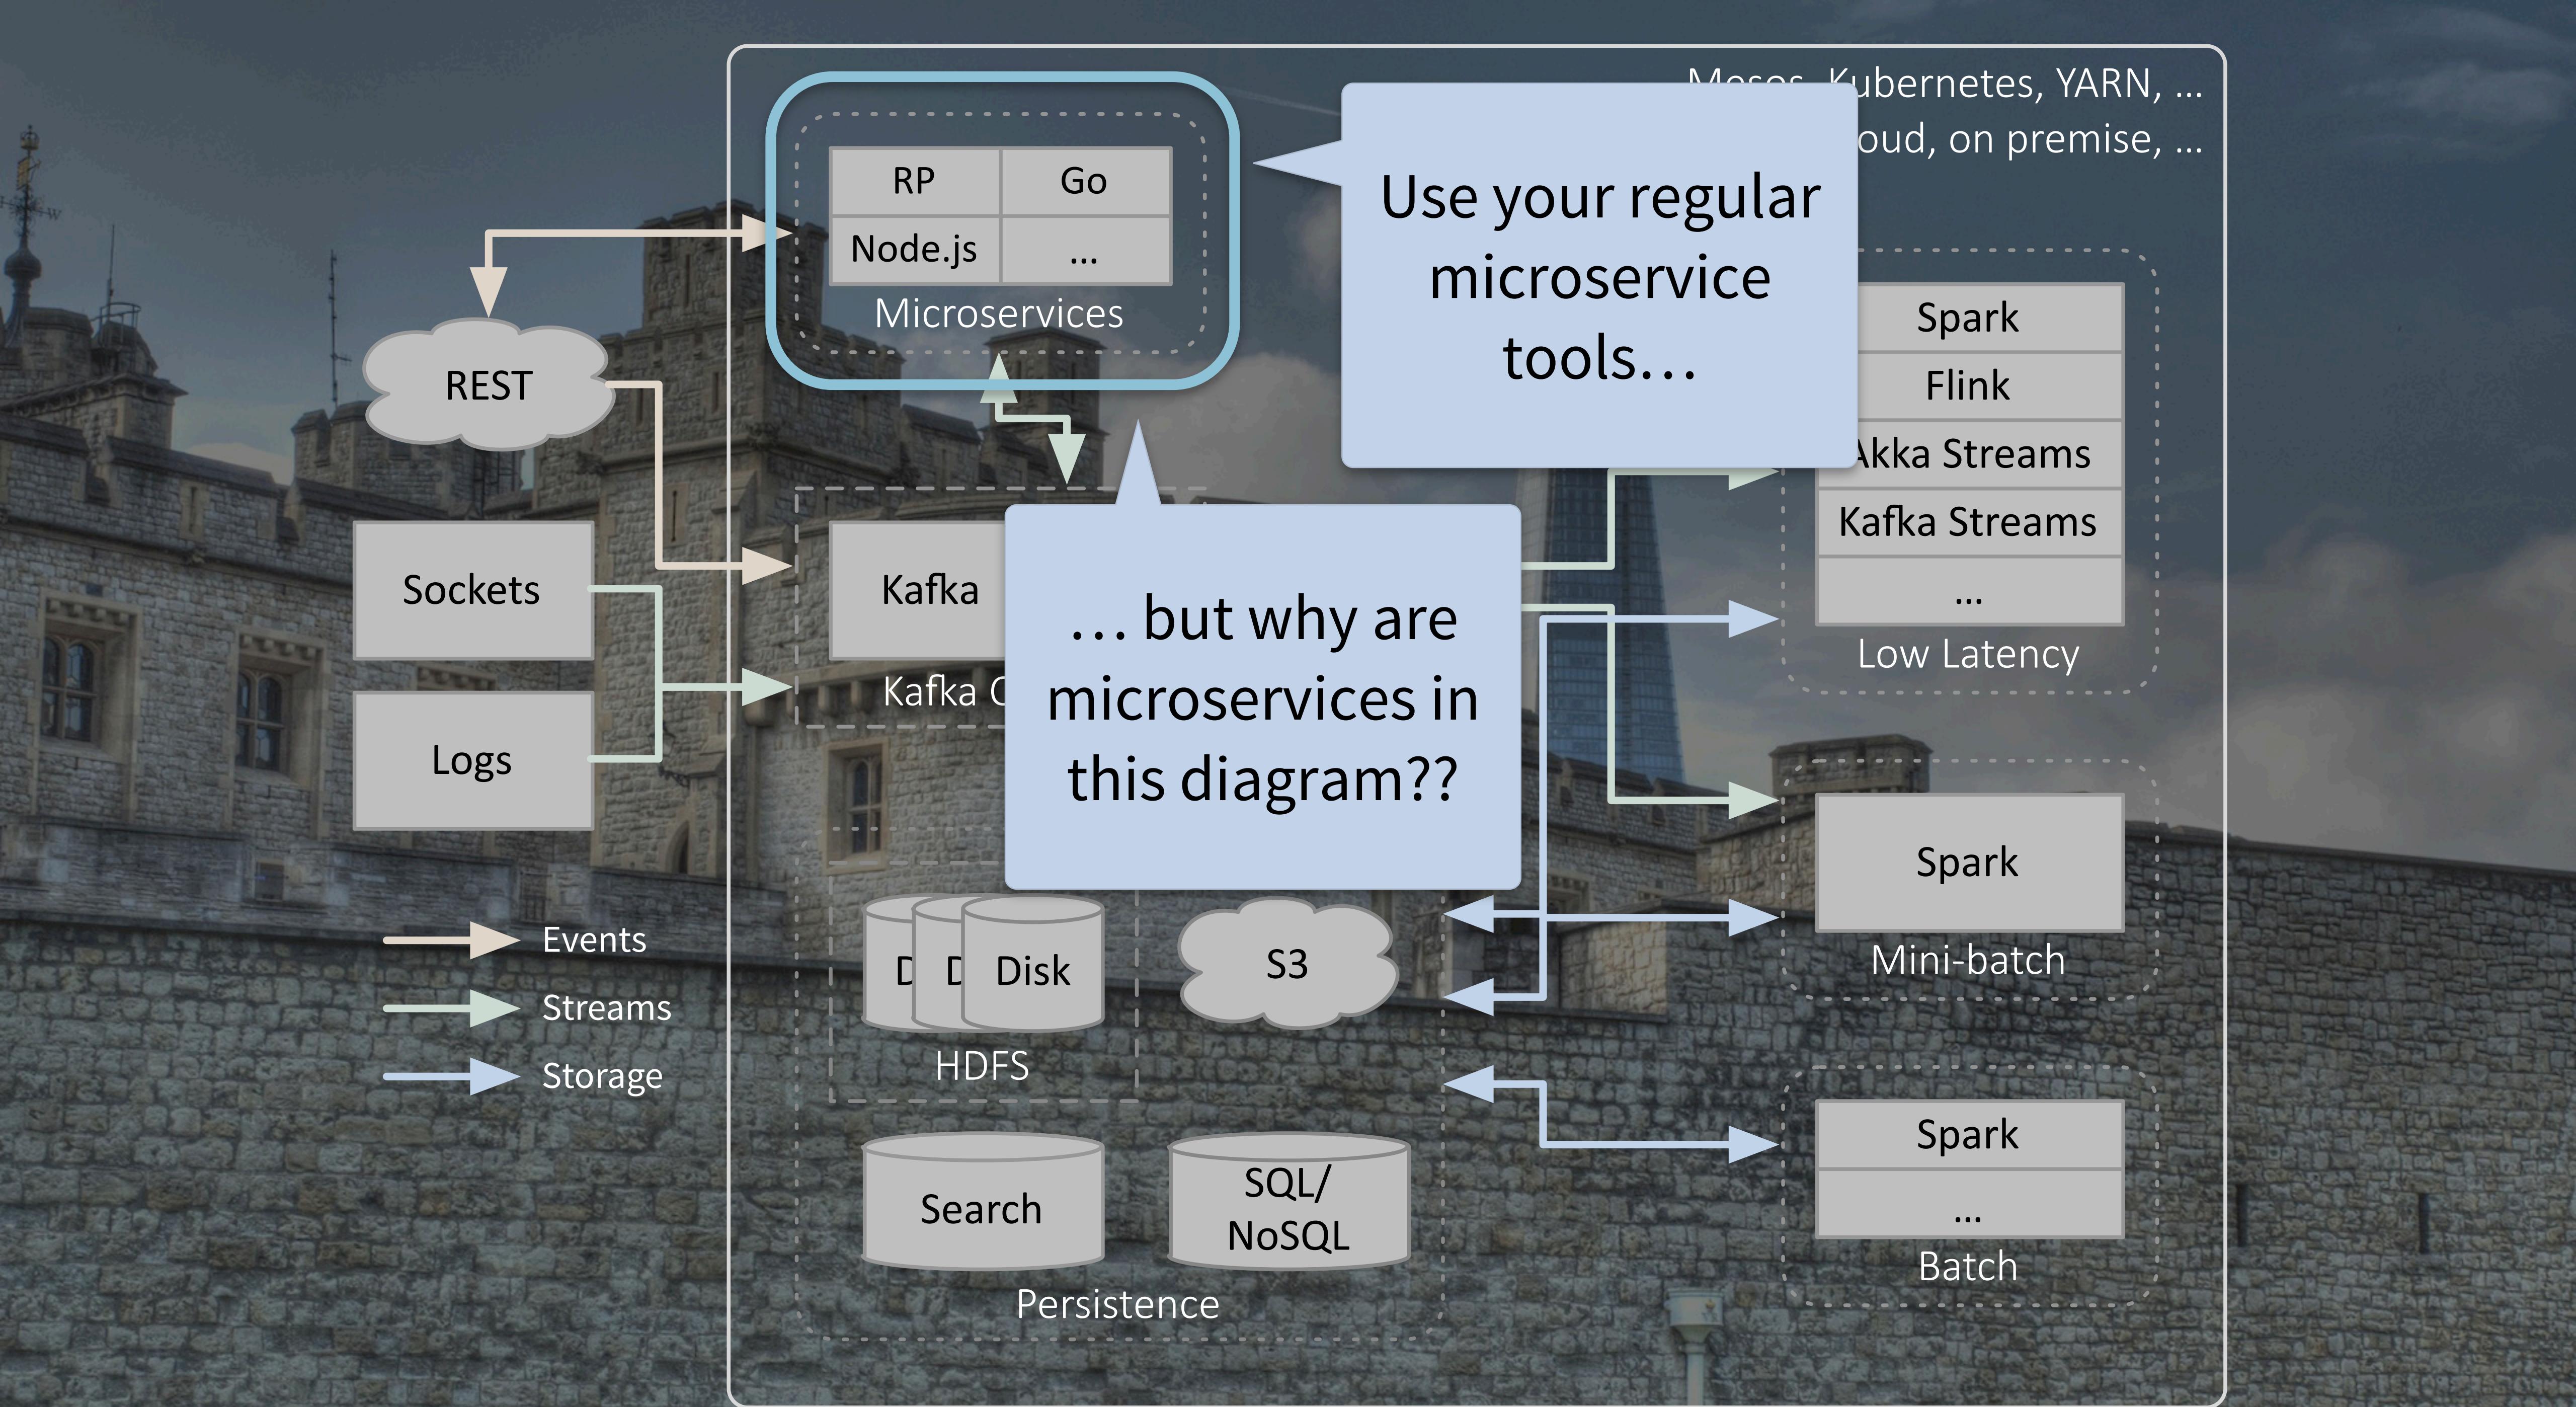
The best known is Spark, which provides rich SQL and dataflow APIs that support batch jobs and streaming jobs. There are two streaming APIs, an older “mini-batch” model, where data is captured in time intervals down to ~100 milliseconds in size, then processed as a small batch. Now there is a new low-latency API.







Data can also be read to and written from almost any storage environment, especially for longer term storage than Kafka is typically used for, and for “non-streaming” applications, like data warehousing, interactive queries, search and indexing, etc.



Microservices in Fast Data

- Run everything in big clusters using Kubernetes or Mesos
 - In the cloud or on-premise



Microservices in Fast Data

- If streaming gives you information faster from new data, you want those answers quickly and easily available in your other services

Microservices in Fast Data

- Streaming raises the bar on data services
 - Compare to batch services, long-running streaming services must be more:
 - Scalable
 - Resilient against failure
 - Flexible

Microservices in Fast Data

- This leads to our last major point...

Organizational Impact



Organizational Impact

- Data engineers have to become good at highly-available microservices
- Microservice engineers have to become good at data

The Recent Past



A Venn diagram consisting of two overlapping circles. The left circle is yellow and contains the word "Services". The right circle is light blue and contains the words "Big Data". The two circles overlap in the center, representing the intersection of the two concepts.

Services

Big Data

Some overlap in concerns, architecture

Until recently, people building “canonical” services for general-purpose IT apps have focused on high availability, scalability, resilience, etc. (i.e., the Reactive principles), recently moving to microservices to do this better. The Big Data world has focused on data storage and cluster scalability, with less need to worry about the Reactive principles. Of course there was some overlap, but they were different spheres...

The Present



Microservices
& Fast Data

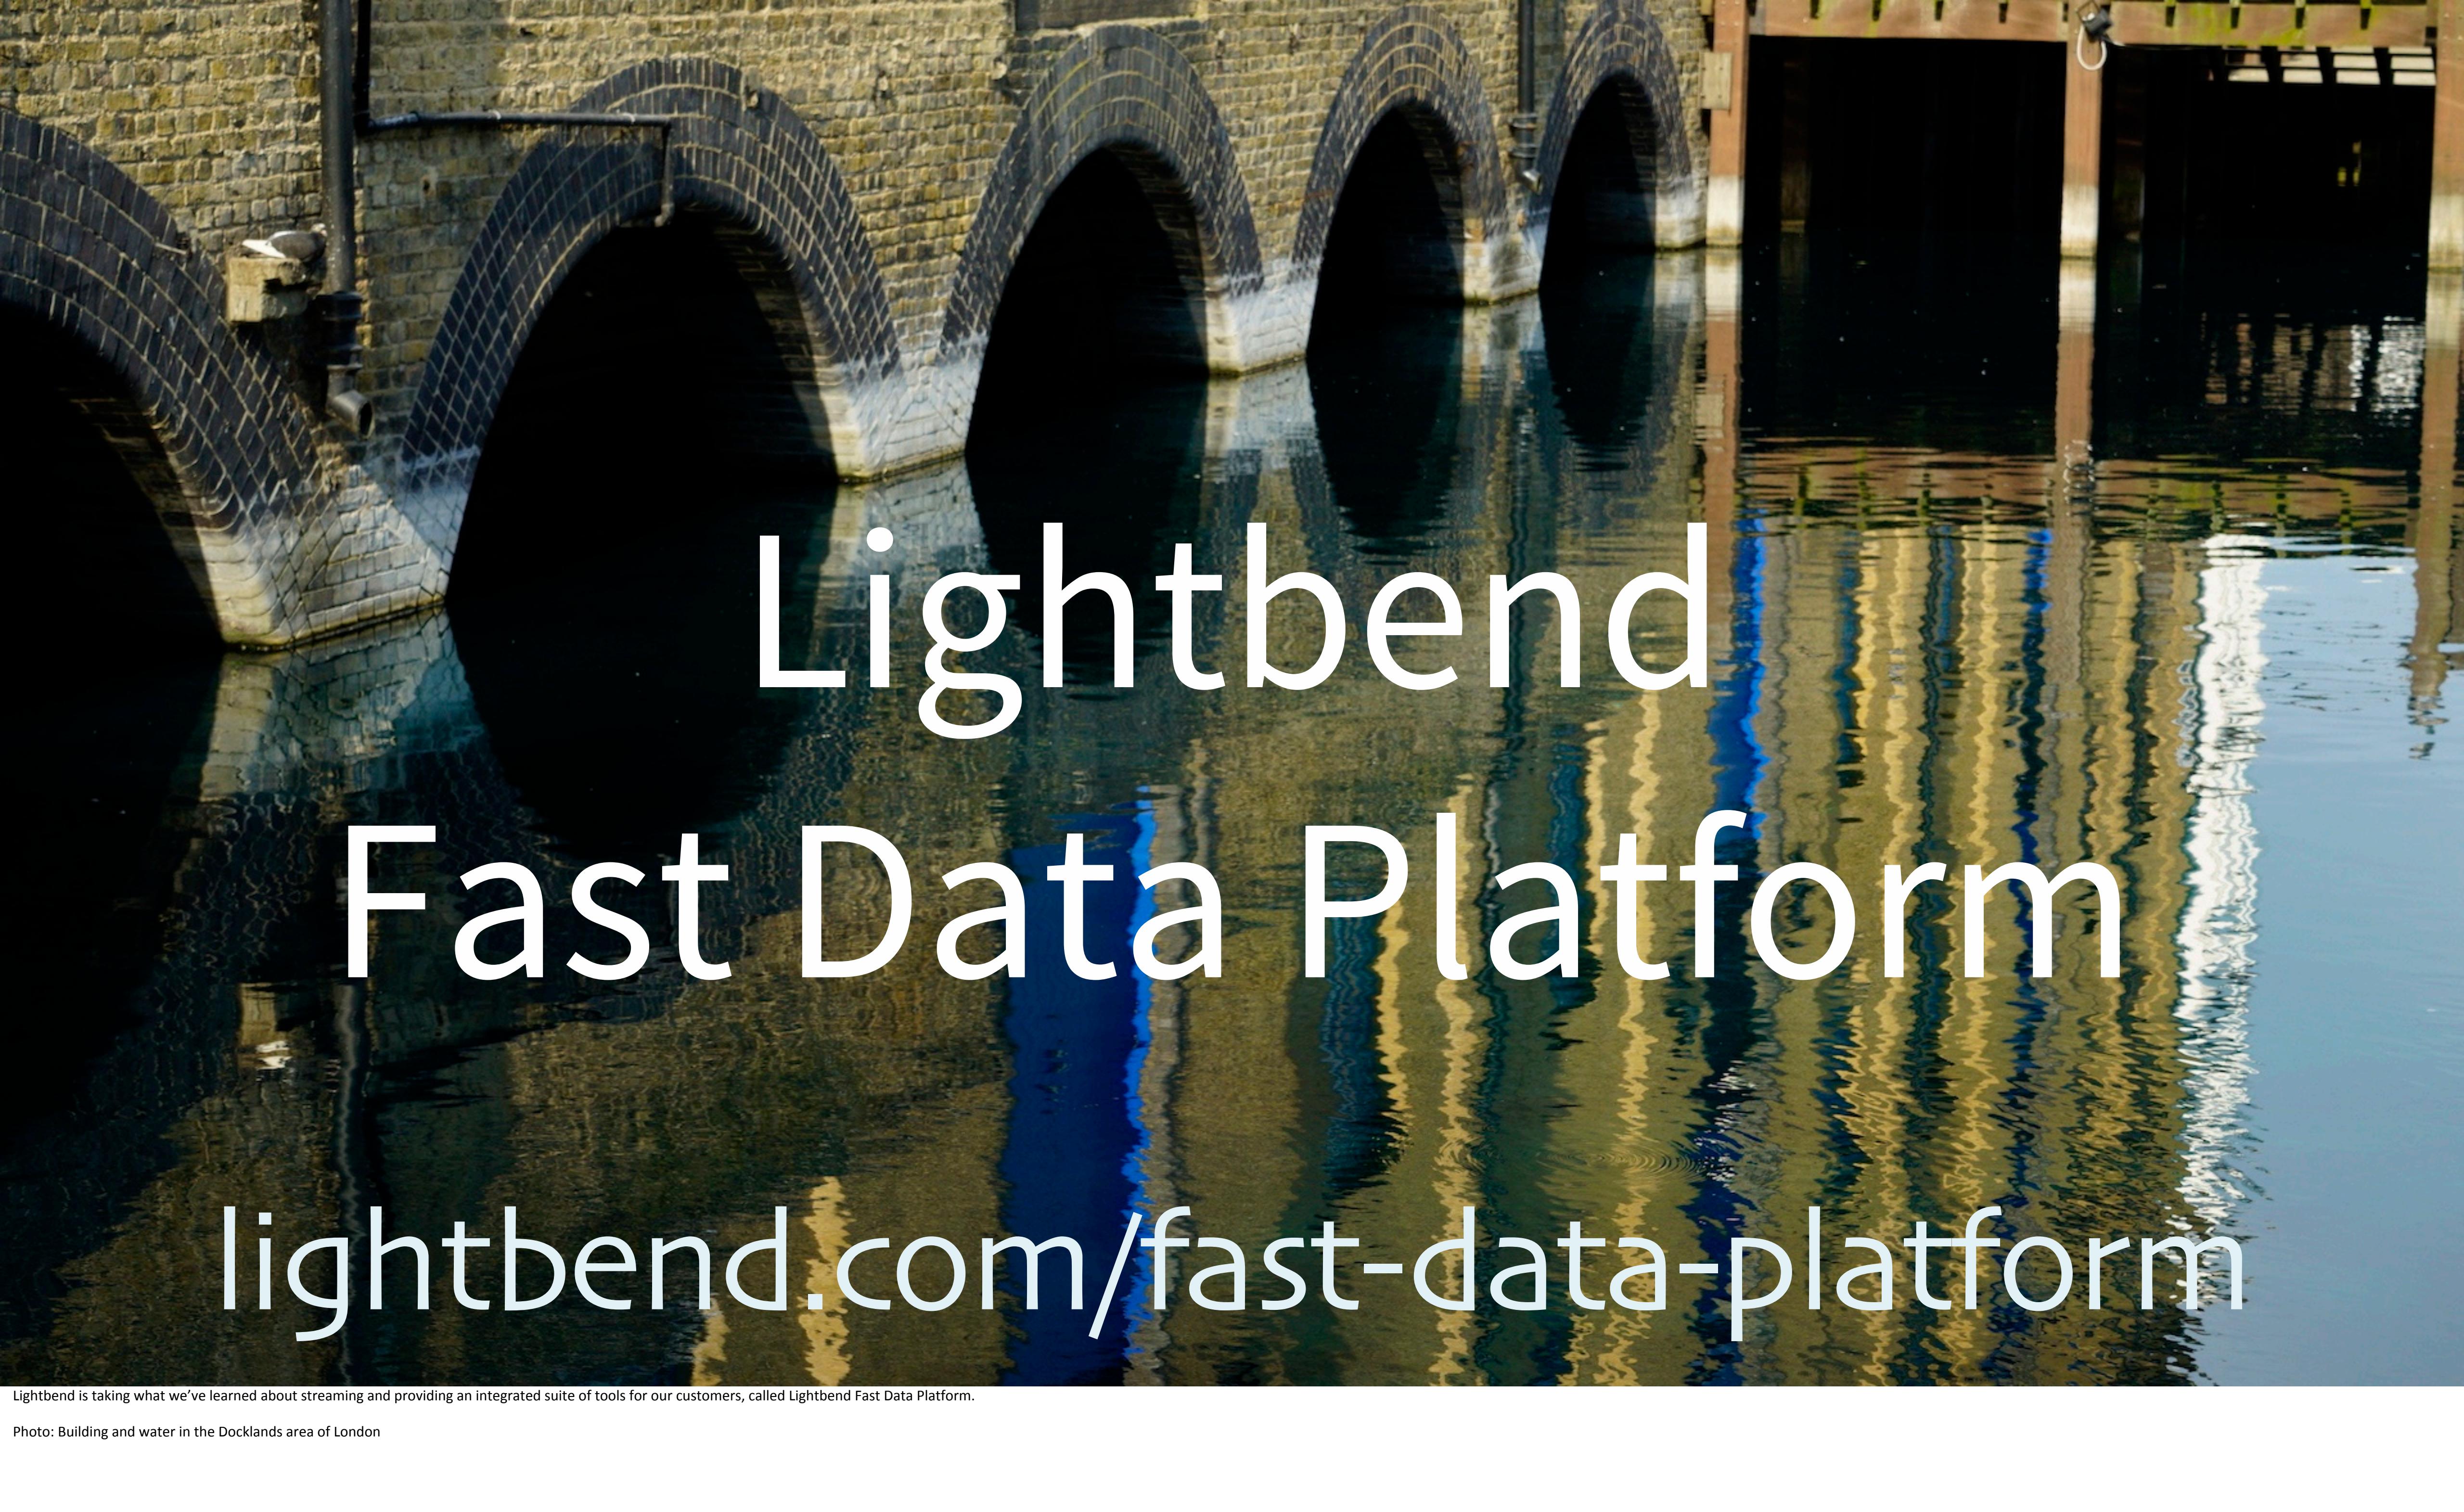
Much more overlap

The Future?



Microservices
for Fast Data

Much more microservice focused

A photograph of a multi-arched brick bridge, likely in a Docklands area, reflected perfectly in the still water below. The bridge's structure is composed of several large, rounded arches made of dark-colored bricks. The reflection in the water creates a symmetrical pattern of arches and shadows. The overall scene is peaceful and architectural.

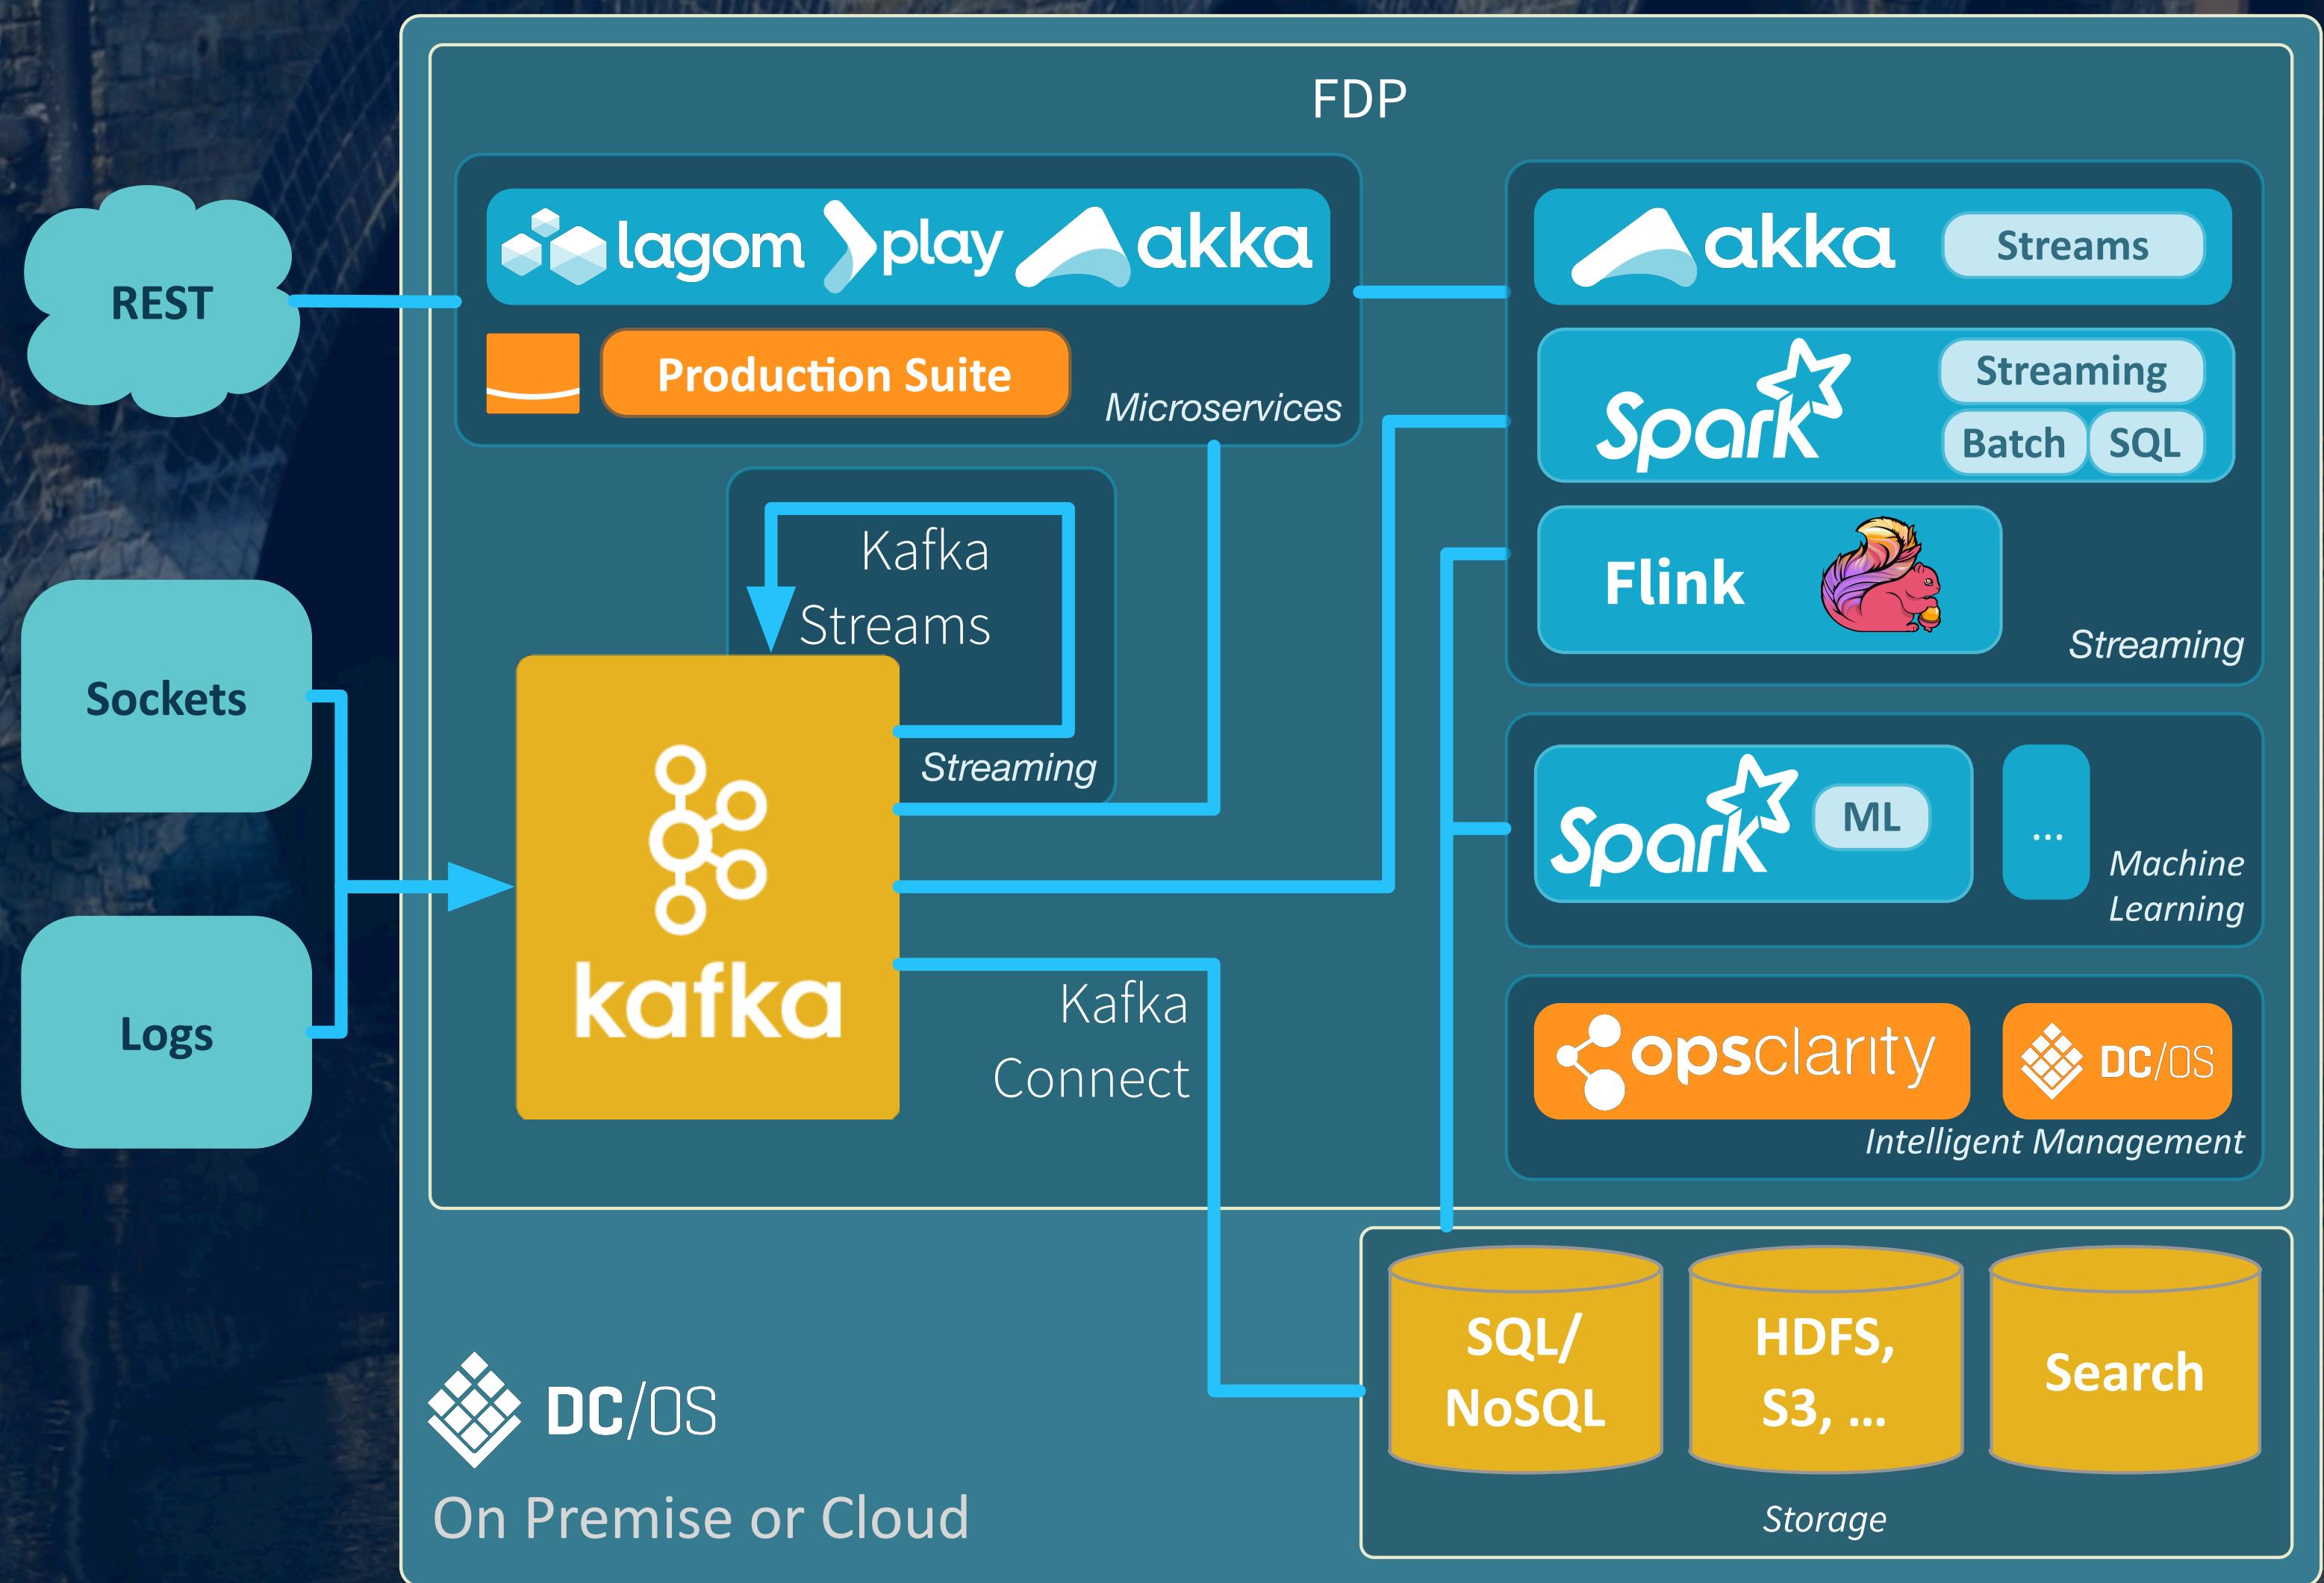
Lightbend Fast Data Platform

lightbend.com/fast-data-platform

Lightbend is taking what we've learned about streaming and providing an integrated suite of tools for our customers, called Lightbend Fast Data Platform.

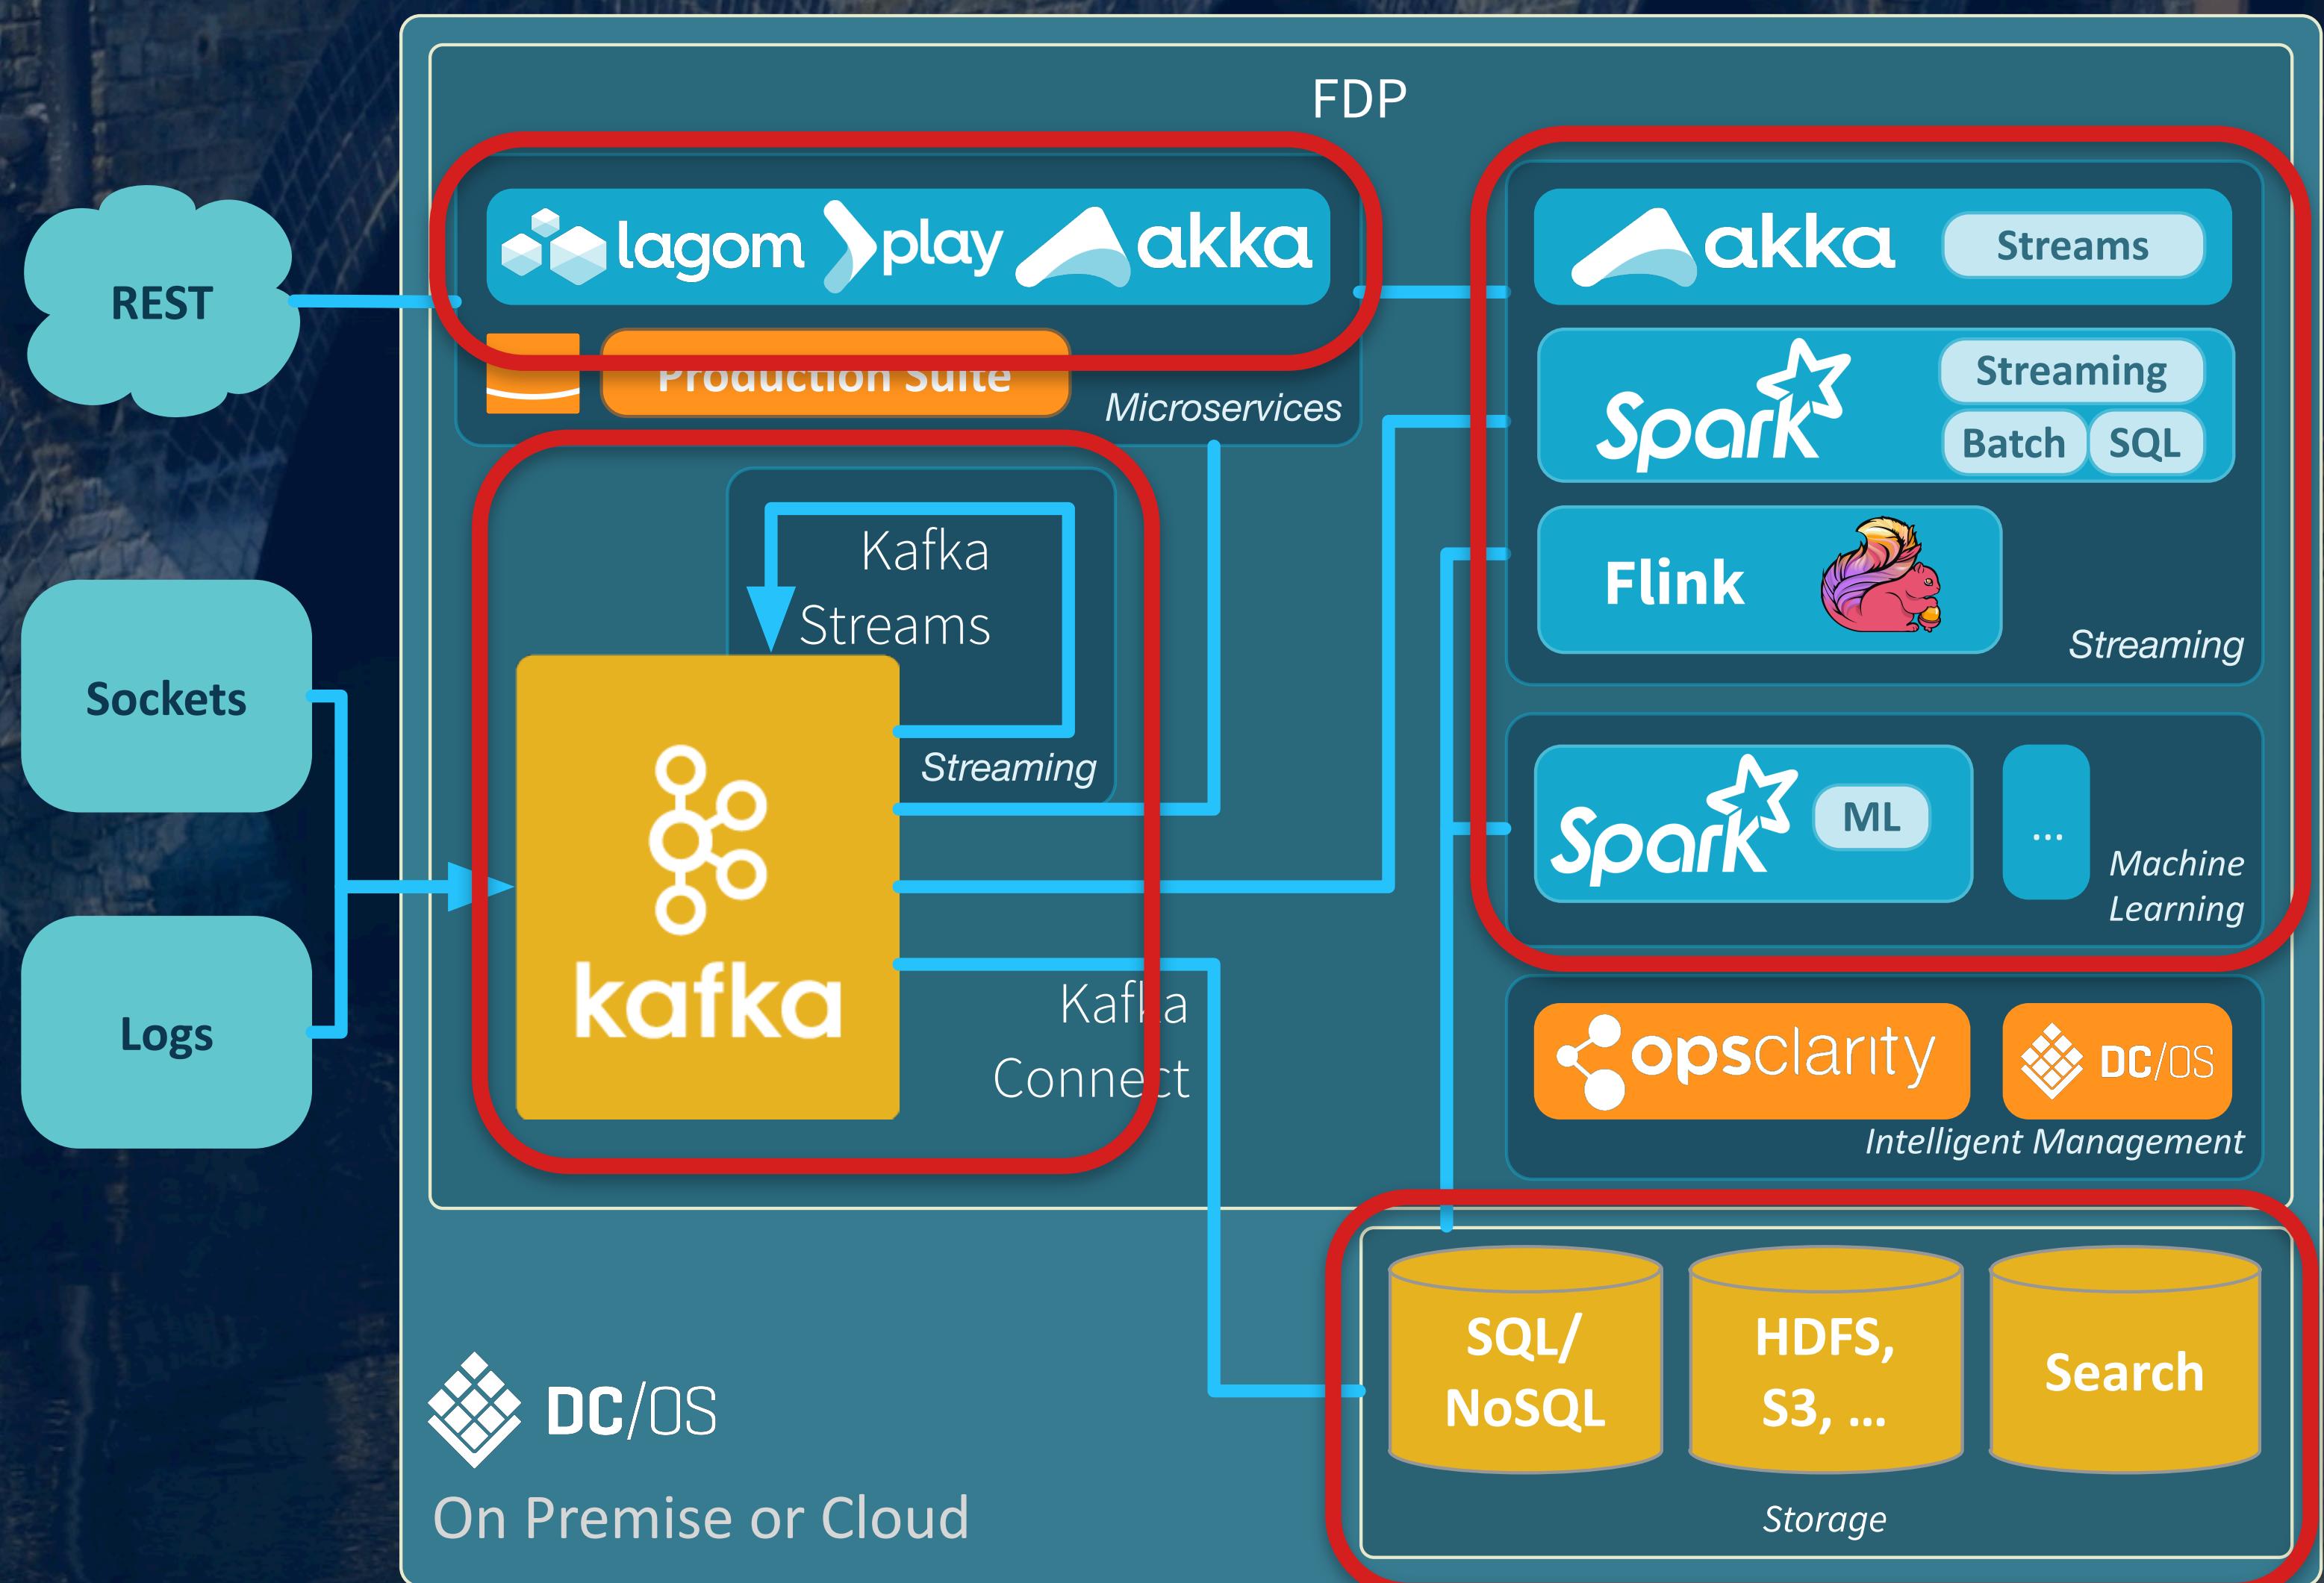
Photo: Building and water in the Docklands area of London

Lightbend Fast Data Platform V1.0



lightbend.com/fast-data-platform

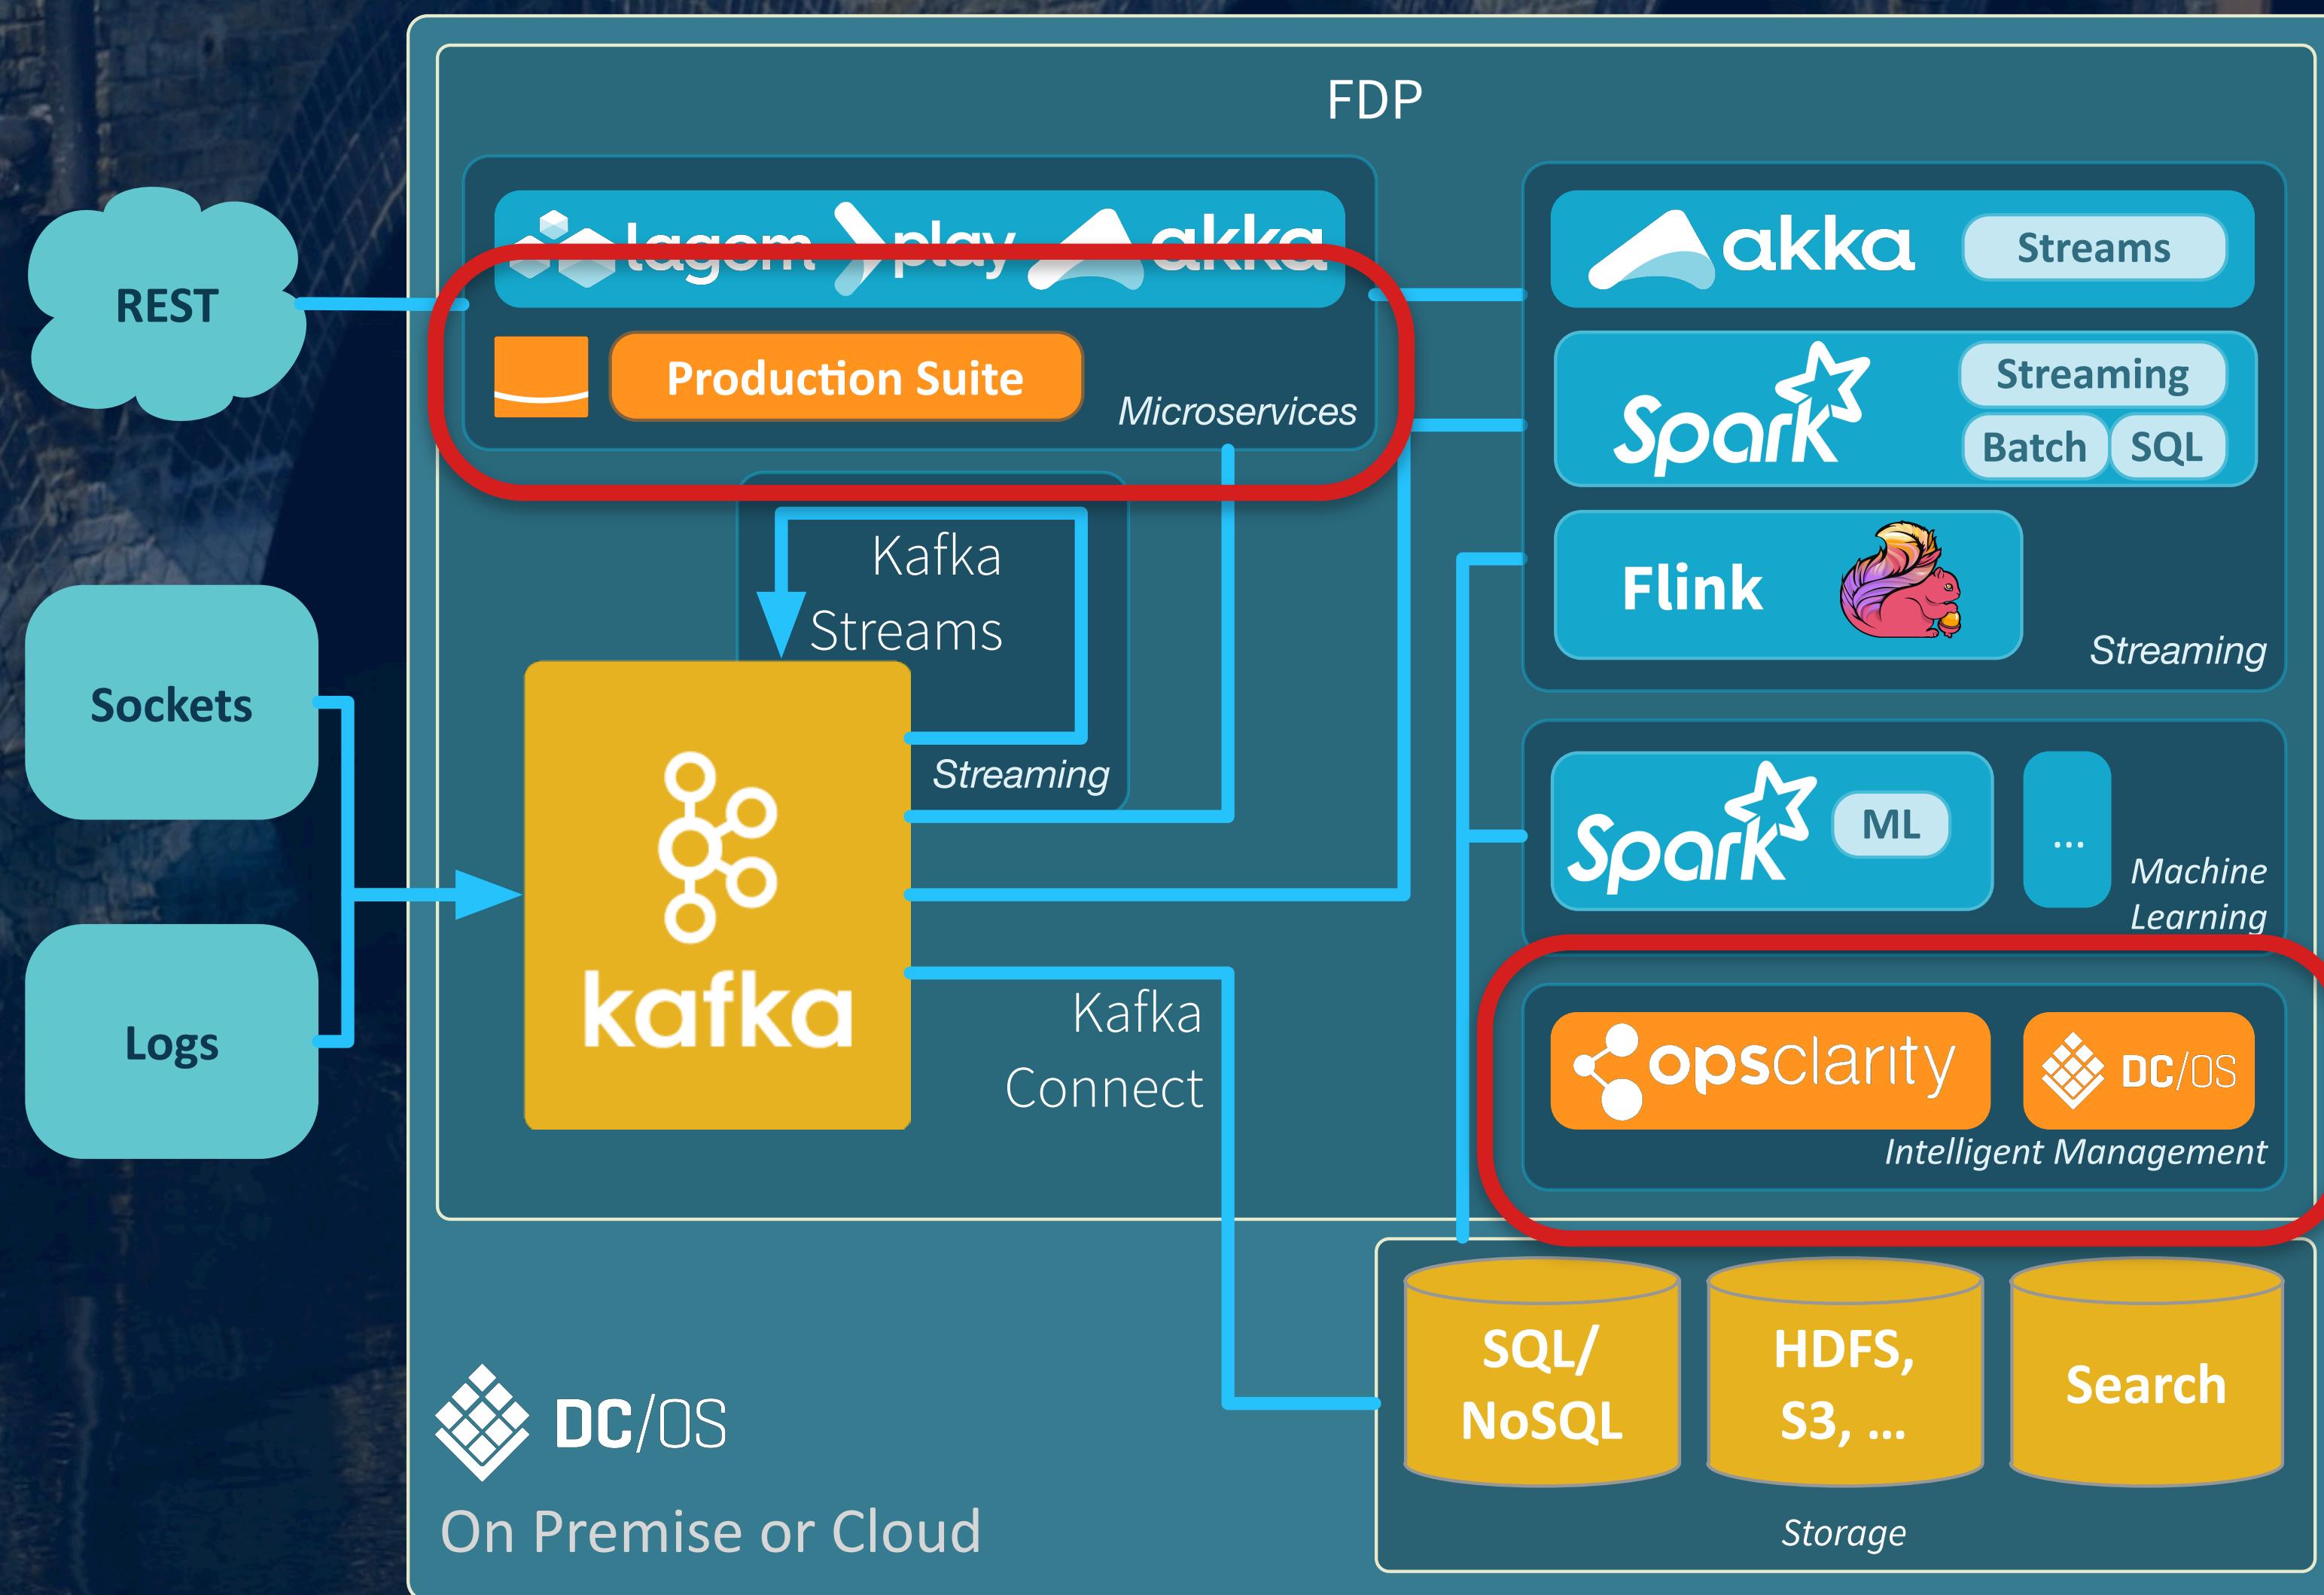
Lightbend Fast Data Platform V1.0



What we
discussed

lightbend.com/fast-data-platform

Lightbend Fast Data Platform V1.0



Plus
management and
monitoring tools

lightbend.com/fast-data-platform

A photograph of the London skyline featuring the London Eye (Millennium Wheel) on the left and the Palace of Westminster with Big Ben on the right. The River Thames flows in the foreground with several boats. The sky is overcast with dramatic clouds.

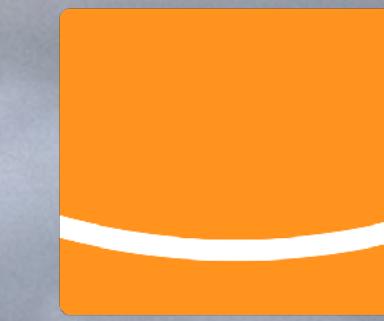
lightbend.com/fast-data-platform

Dean Wampler, Ph.D.

dean@lightbend.com

@deanwampler

polyglotprogramming.com/talks



Lightbend