

NoSQL Search RoadShow  
San Francisco, June 6, 2013  
[dean@concurrentthought.com](mailto:dean@concurrentthought.com)  
[@deanwampler](https://twitter.com/deanwampler)  
[polyglotprogramming.com/talks](http://polyglotprogramming.com/talks)



# From Big Data to Big Information

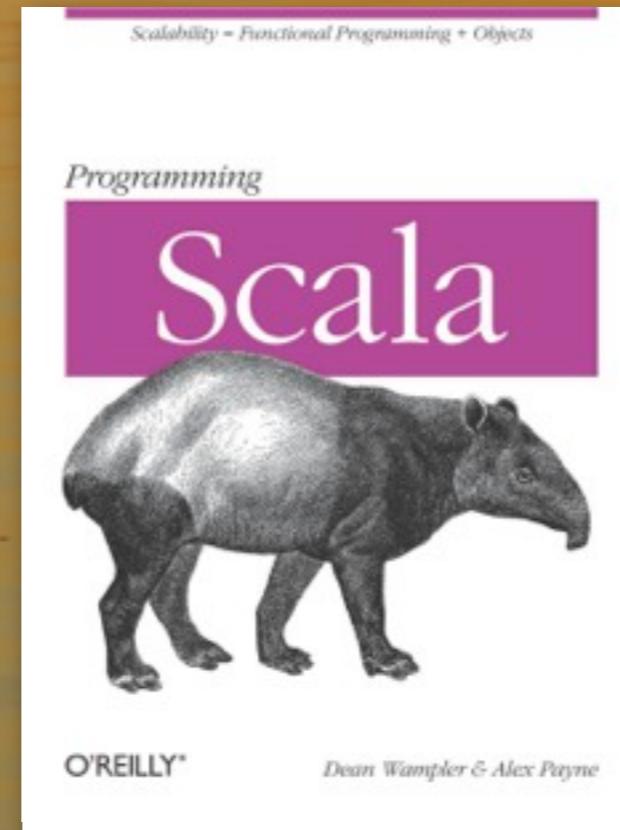
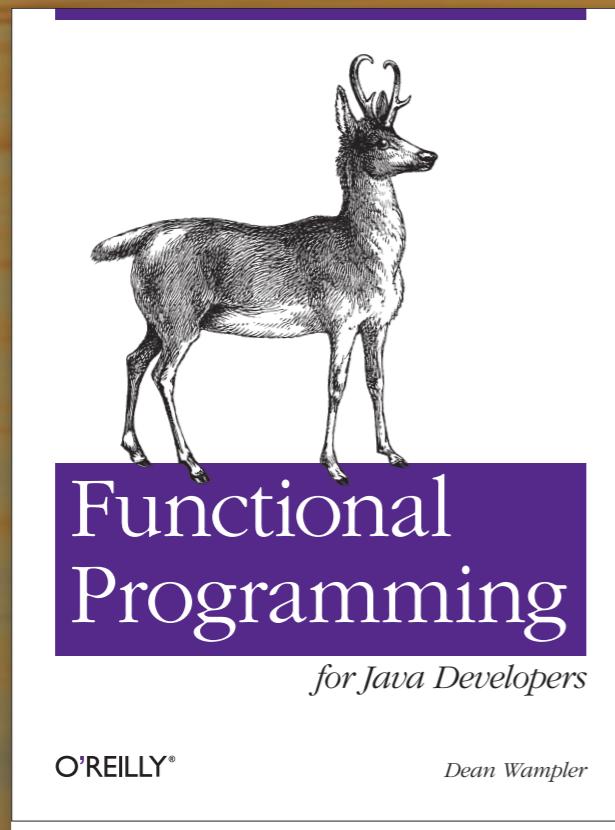


Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Copyright © Dean Wampler, 2011–2013, All Rights Reserved. Photos can only be used with permission. Otherwise, the content is free to use.  
Photo: San Francisco Bay, just south of the airport in Burlingame, before sunrise.

# Who am I?



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

My books...

Photo: San Francisco Bay, Burlingame, around sunrise.

# Why this talk?

*Hadoop hype cycle:  
But what gives us  
actual value?*

Sunday, June 9, 13

This talk reflects my experiences working on “Big Data” projects, mostly using Hadoop, with many clients. People sometimes jump on this bandwagon to avoid “being left behind” or to boost their career. Sometimes, it doesn’t make sense for their actual needs. Big Data itself doesn’t do you any good. It’s the information we extract that’s important, so let’s see how different technologies address specific problems.

Photo: Photo: San Francisco Bay, Burlingame, around sunrise.

# What Is Big Data?



**DevOps Borat** @DEVOPS\_BORAT

8 Jan

Big Data is any thing which is crash Excel.

[Expand](#)



**DevOps Borat** @DEVOPS\_BORAT

6 Feb

Small Data is when is fit in RAM. Big Data is when is crash because is not fit in RAM.

[Expand](#)

# Big Data

Data so big that traditional solutions are too slow, too small, or too expensive to use.



Hat tip: Bob Korbus

# 3 Trends



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

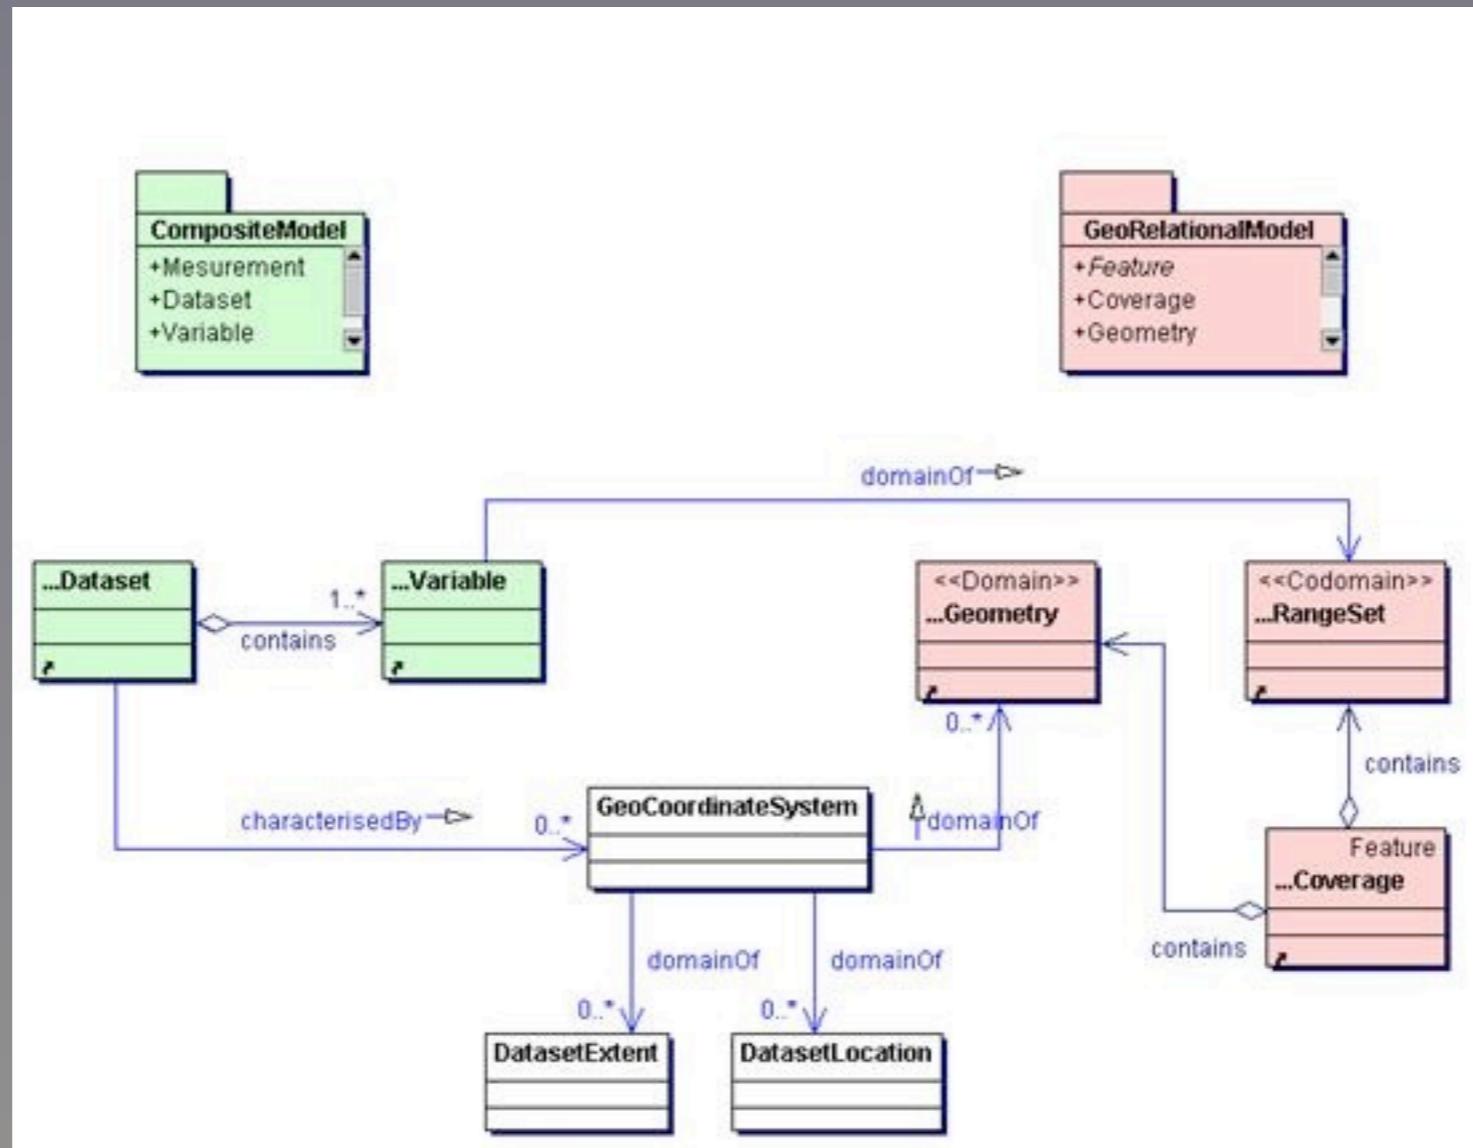
Three trends to organizer our thinking...

Photo: Gull on a pier near Fort Mason, SF

# Data Size ↑



# Formal Schemas ↓



There is less emphasis on “formal” schemas and domain models, i.e., both relational models of data and OO models, because data schemas and sources change rapidly, and we need to integrate so many disparate sources of data. So, using relatively-agnostic software, e.g., collections of things where the software is more agnostic about the structure of the data and the domain, tends to be faster to develop, test, and deploy. Put another way, we find it more useful to build somewhat agnostic applications and drive their behavior through data...

# Data-Driven Programs ↑



Sunday, June 9, 13

This is the 2nd generation “Stanley”, the most successful self-driving car ever built (by a Google-Stanford) team. Machine learning is growing in importance. Here, generic algorithms and data structures are trained to represent the “world” using data, rather than encoding a model of the world in the software itself. It’s another example of generic algorithms that produce the desired behavior by being application agnostic and data driven, rather than hard-coding a model of the world. (In practice, however, a balance is struck between completely agnostic apps and some engineering towards for the specific problem, as you might expect...)

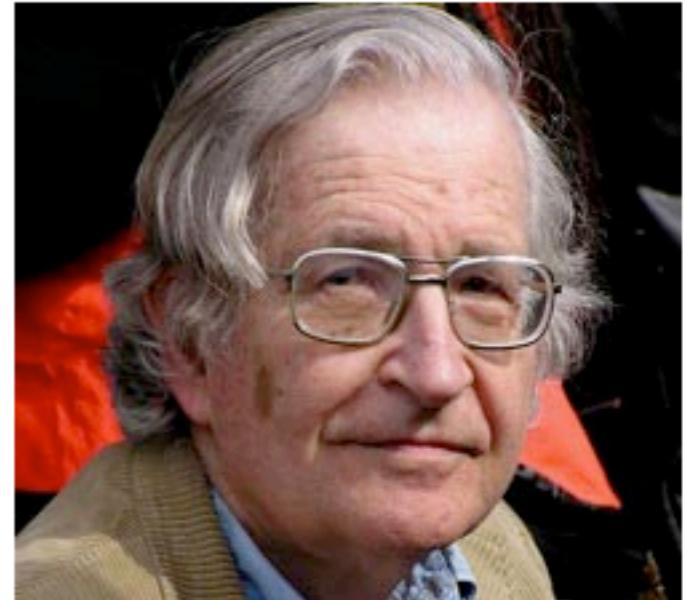
# Probabilistic Models vs. Formal Grammars

[tor.com/blogs/...](http://www.tor.com/blogs/...)

## Norvig vs. Chomsky and the Fight for the Future of AI

KEVIN GOLD

When the Director of Research for Google compares one of the most highly regarded linguists of all time to Bill O'Reilly, you know it is *on*. Recently, Peter Norvig, Google's Director of Research and co-author of [the most popular artificial intelligence textbook in the world](#), wrote a [webpage](#) extensively criticizing Noam Chomsky, arguably the most influential linguist in the world. Their disagreement points to a revolution in artificial intelligence that, like many revolutions, threatens to destroy as much as it improves. Chomsky, one of the old guard, wishes for an elegant theory of intelligence and language that looks past human fallibility to try to see simple structure underneath. Norvig, meanwhile, represents the new philosophy: truth by statistics,



Chomsky photo by Duncan Rawlinson and his Online Photography School. Norvig photo by Peter Norvig

# Big Data Architectures

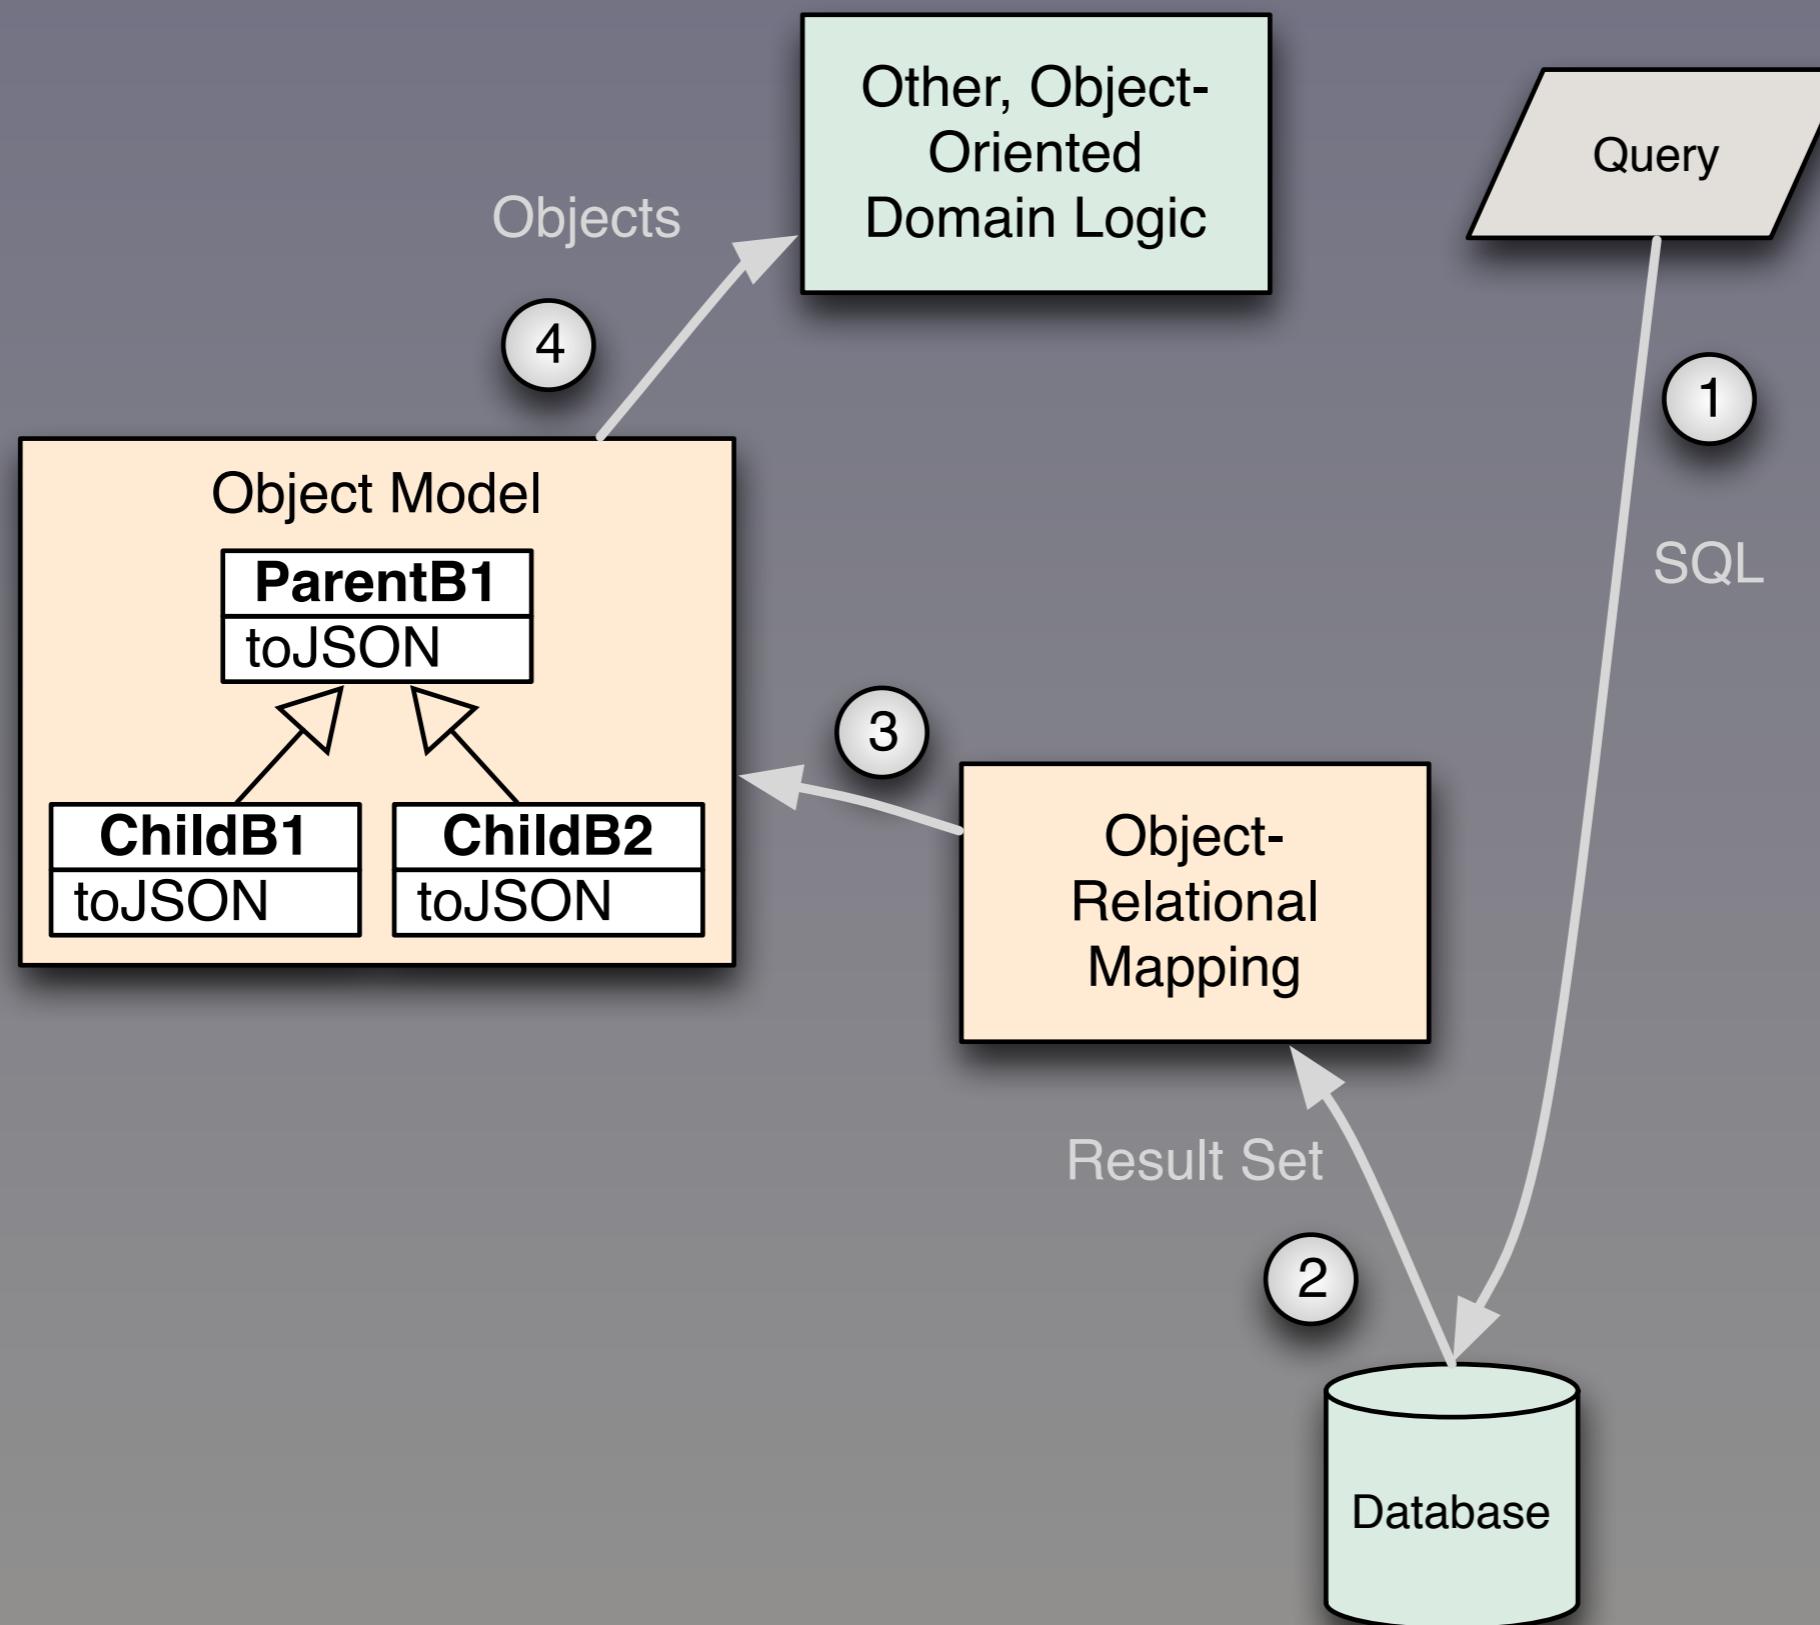


Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

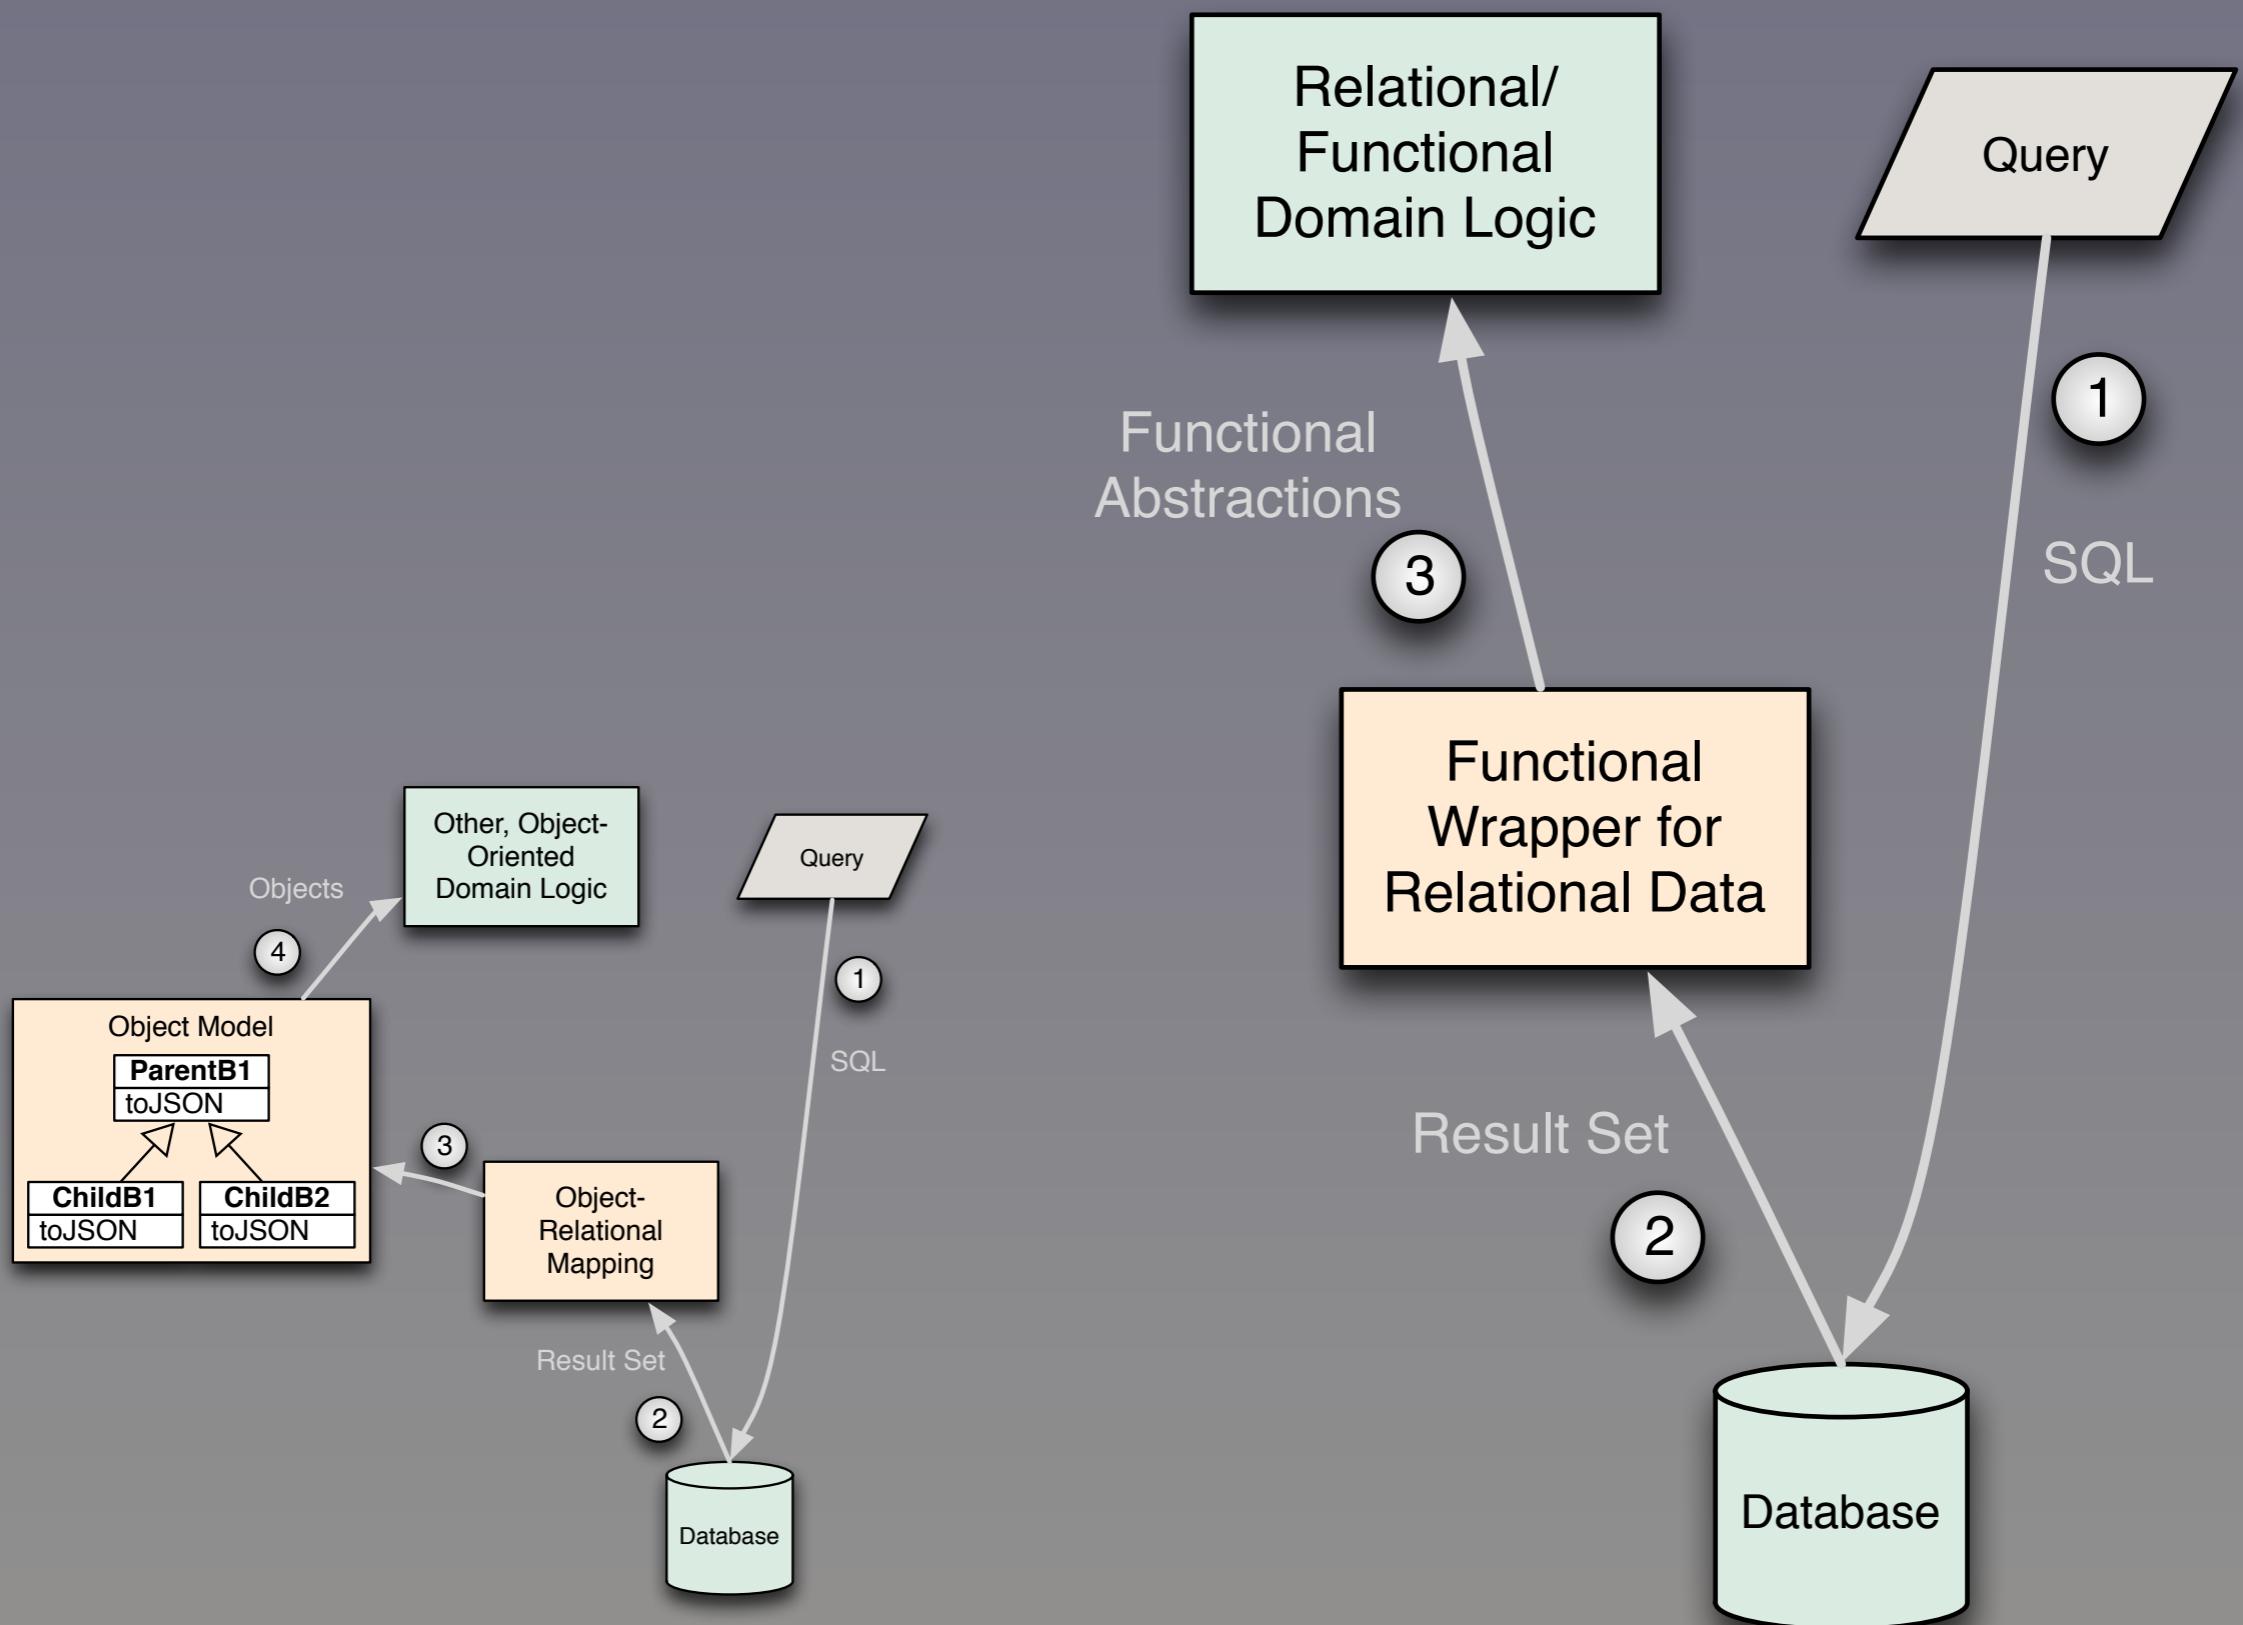
What should software architectures look like for these kinds of systems?

Photo: Light on the Boardwalk and Coit Tower.

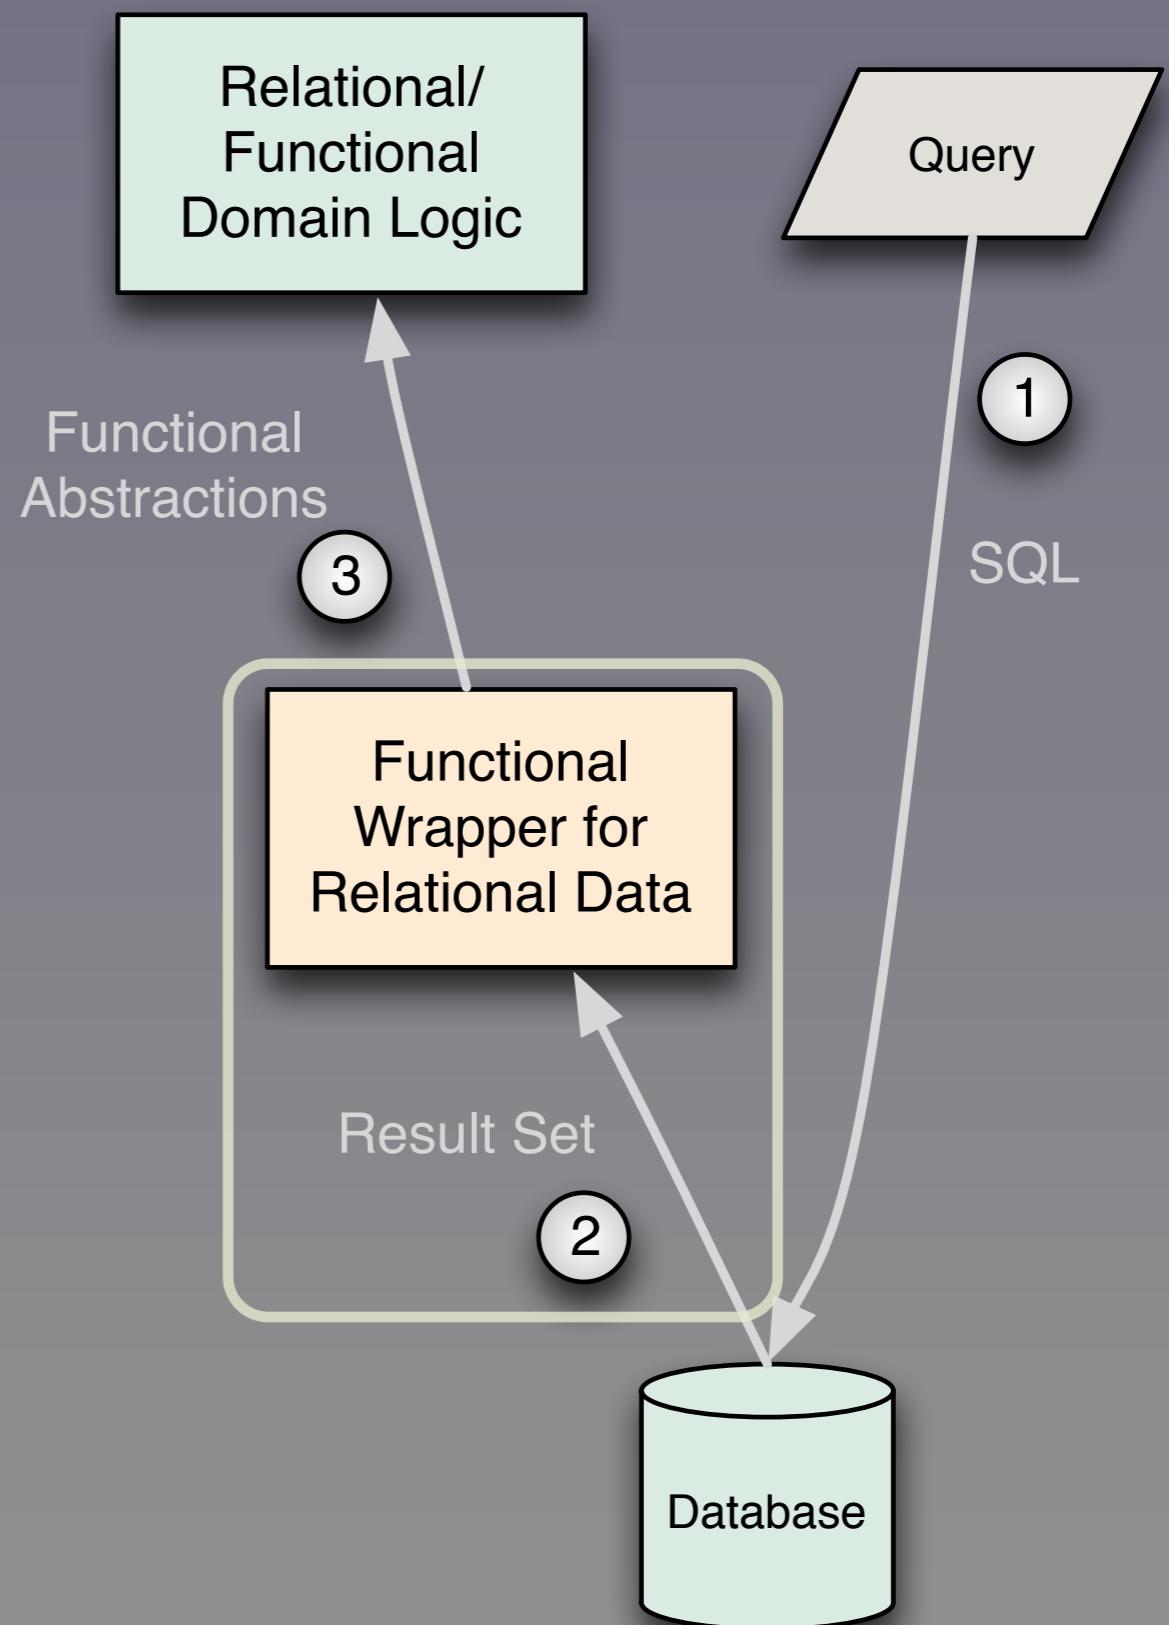


Traditionally, we've kept a rich, in-memory domain model requiring an ORM to convert persistent data into the model. This is resource overhead and complexity we can't afford in big data systems. Rather, we should treat the result set as it is, a particular kind of collection, do the minimal transformation required to exploit our collections libraries and classes representing some domain concepts (e.g., Address, StockOption, etc.), then write functional code to implement business logic (or drive emergent behavior with machine learning algorithms...).

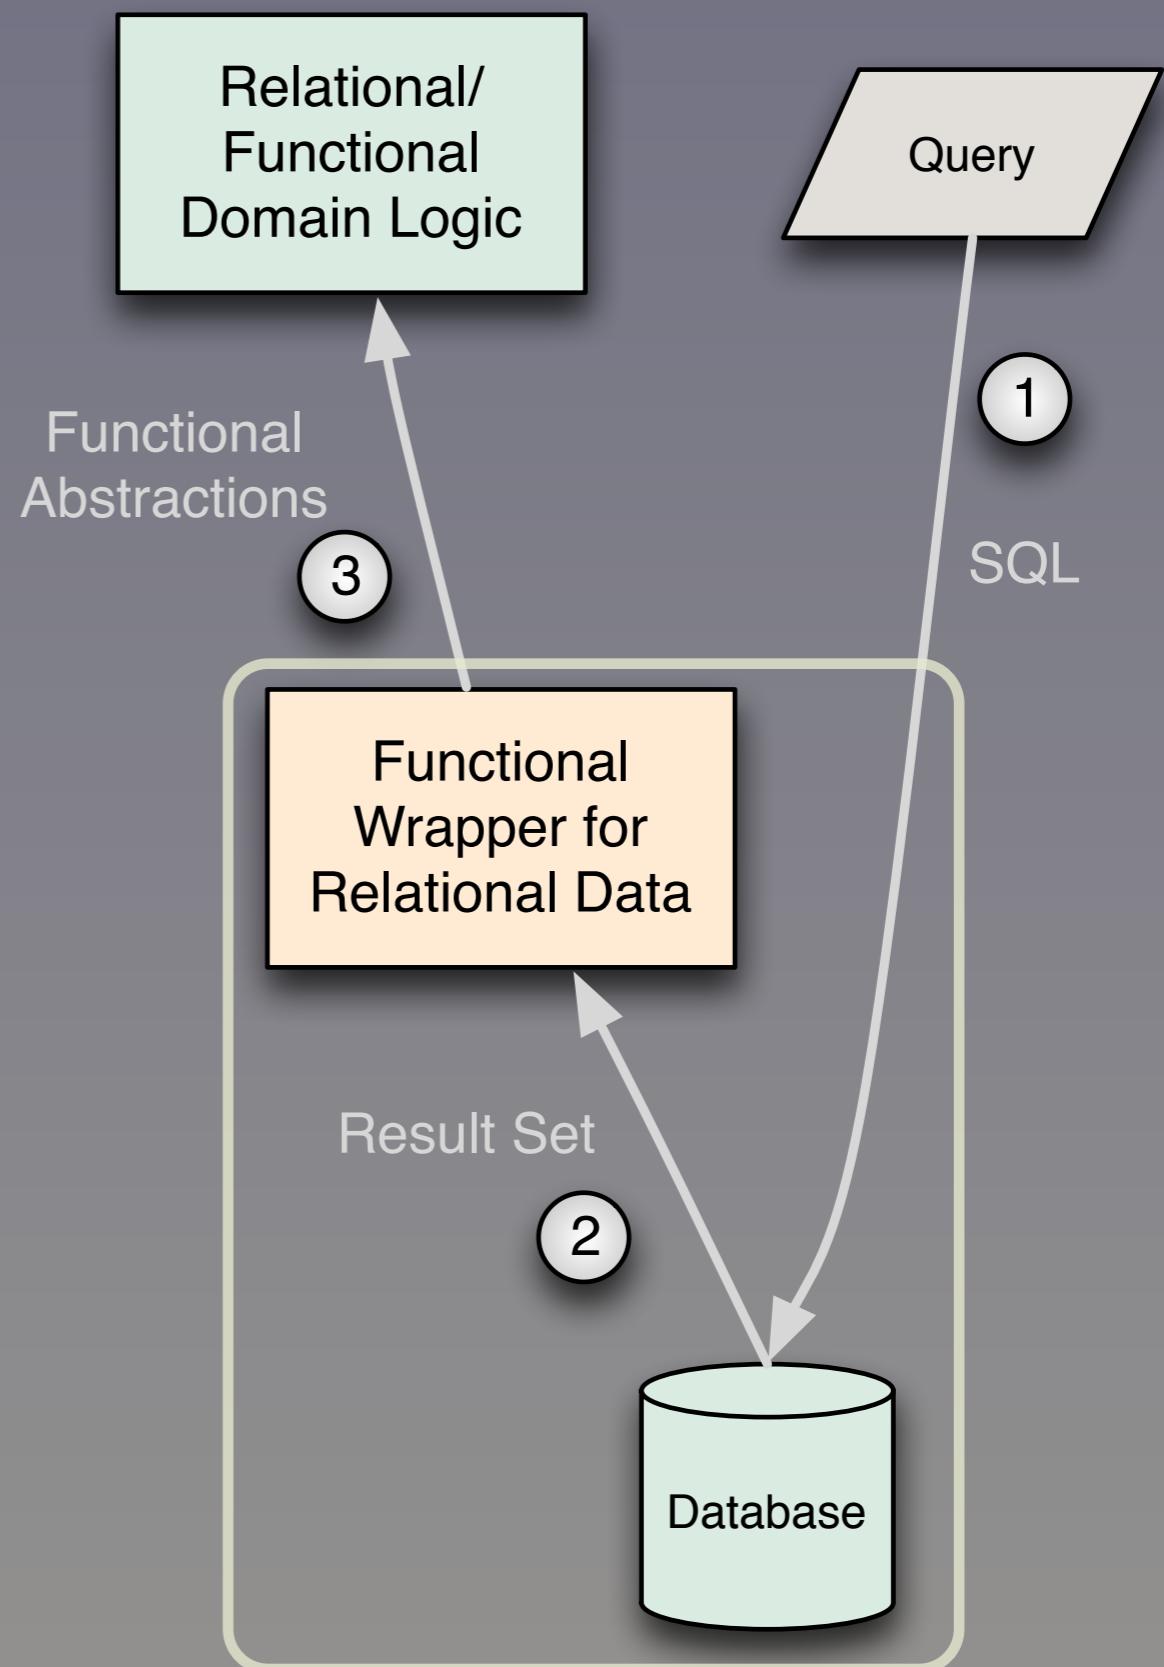
The `toJSON` methods are there because we often convert these object graphs back into fundamental structures, such as the maps and arrays of JSON so we can send them to the browser!

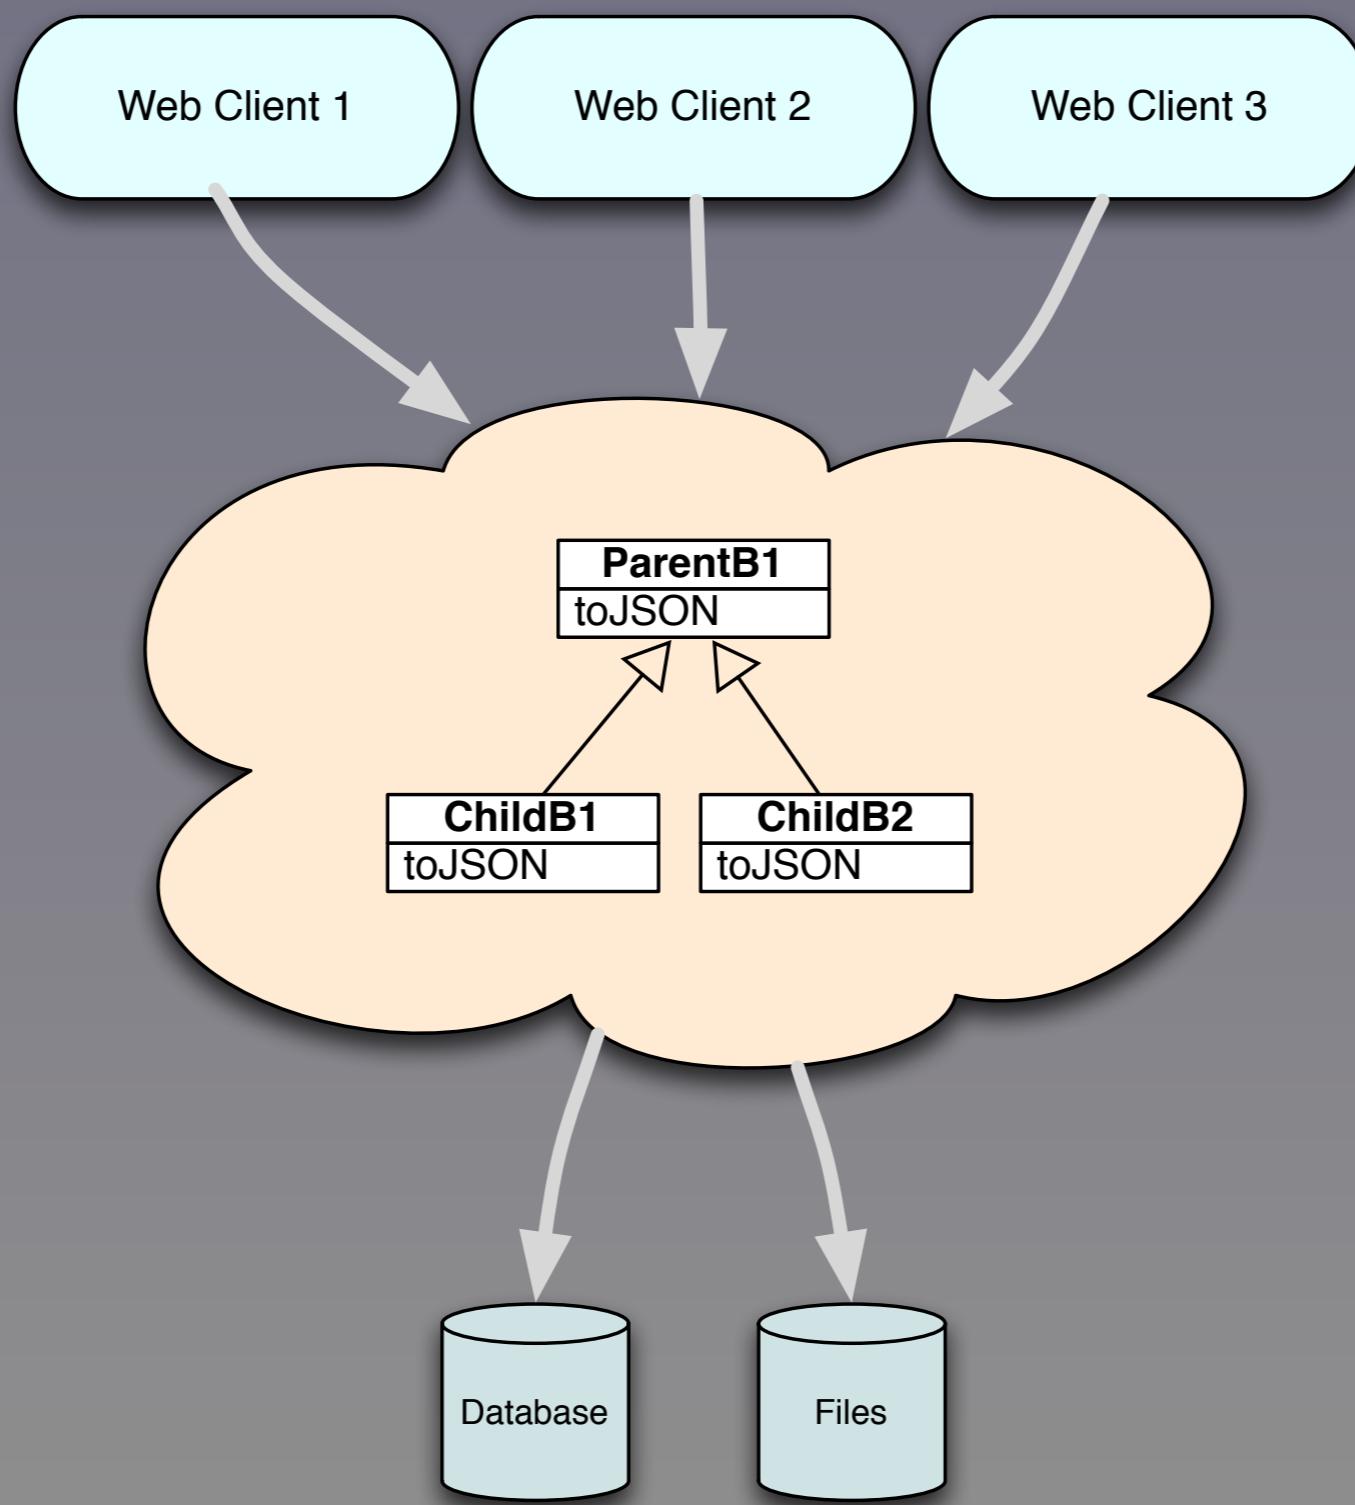


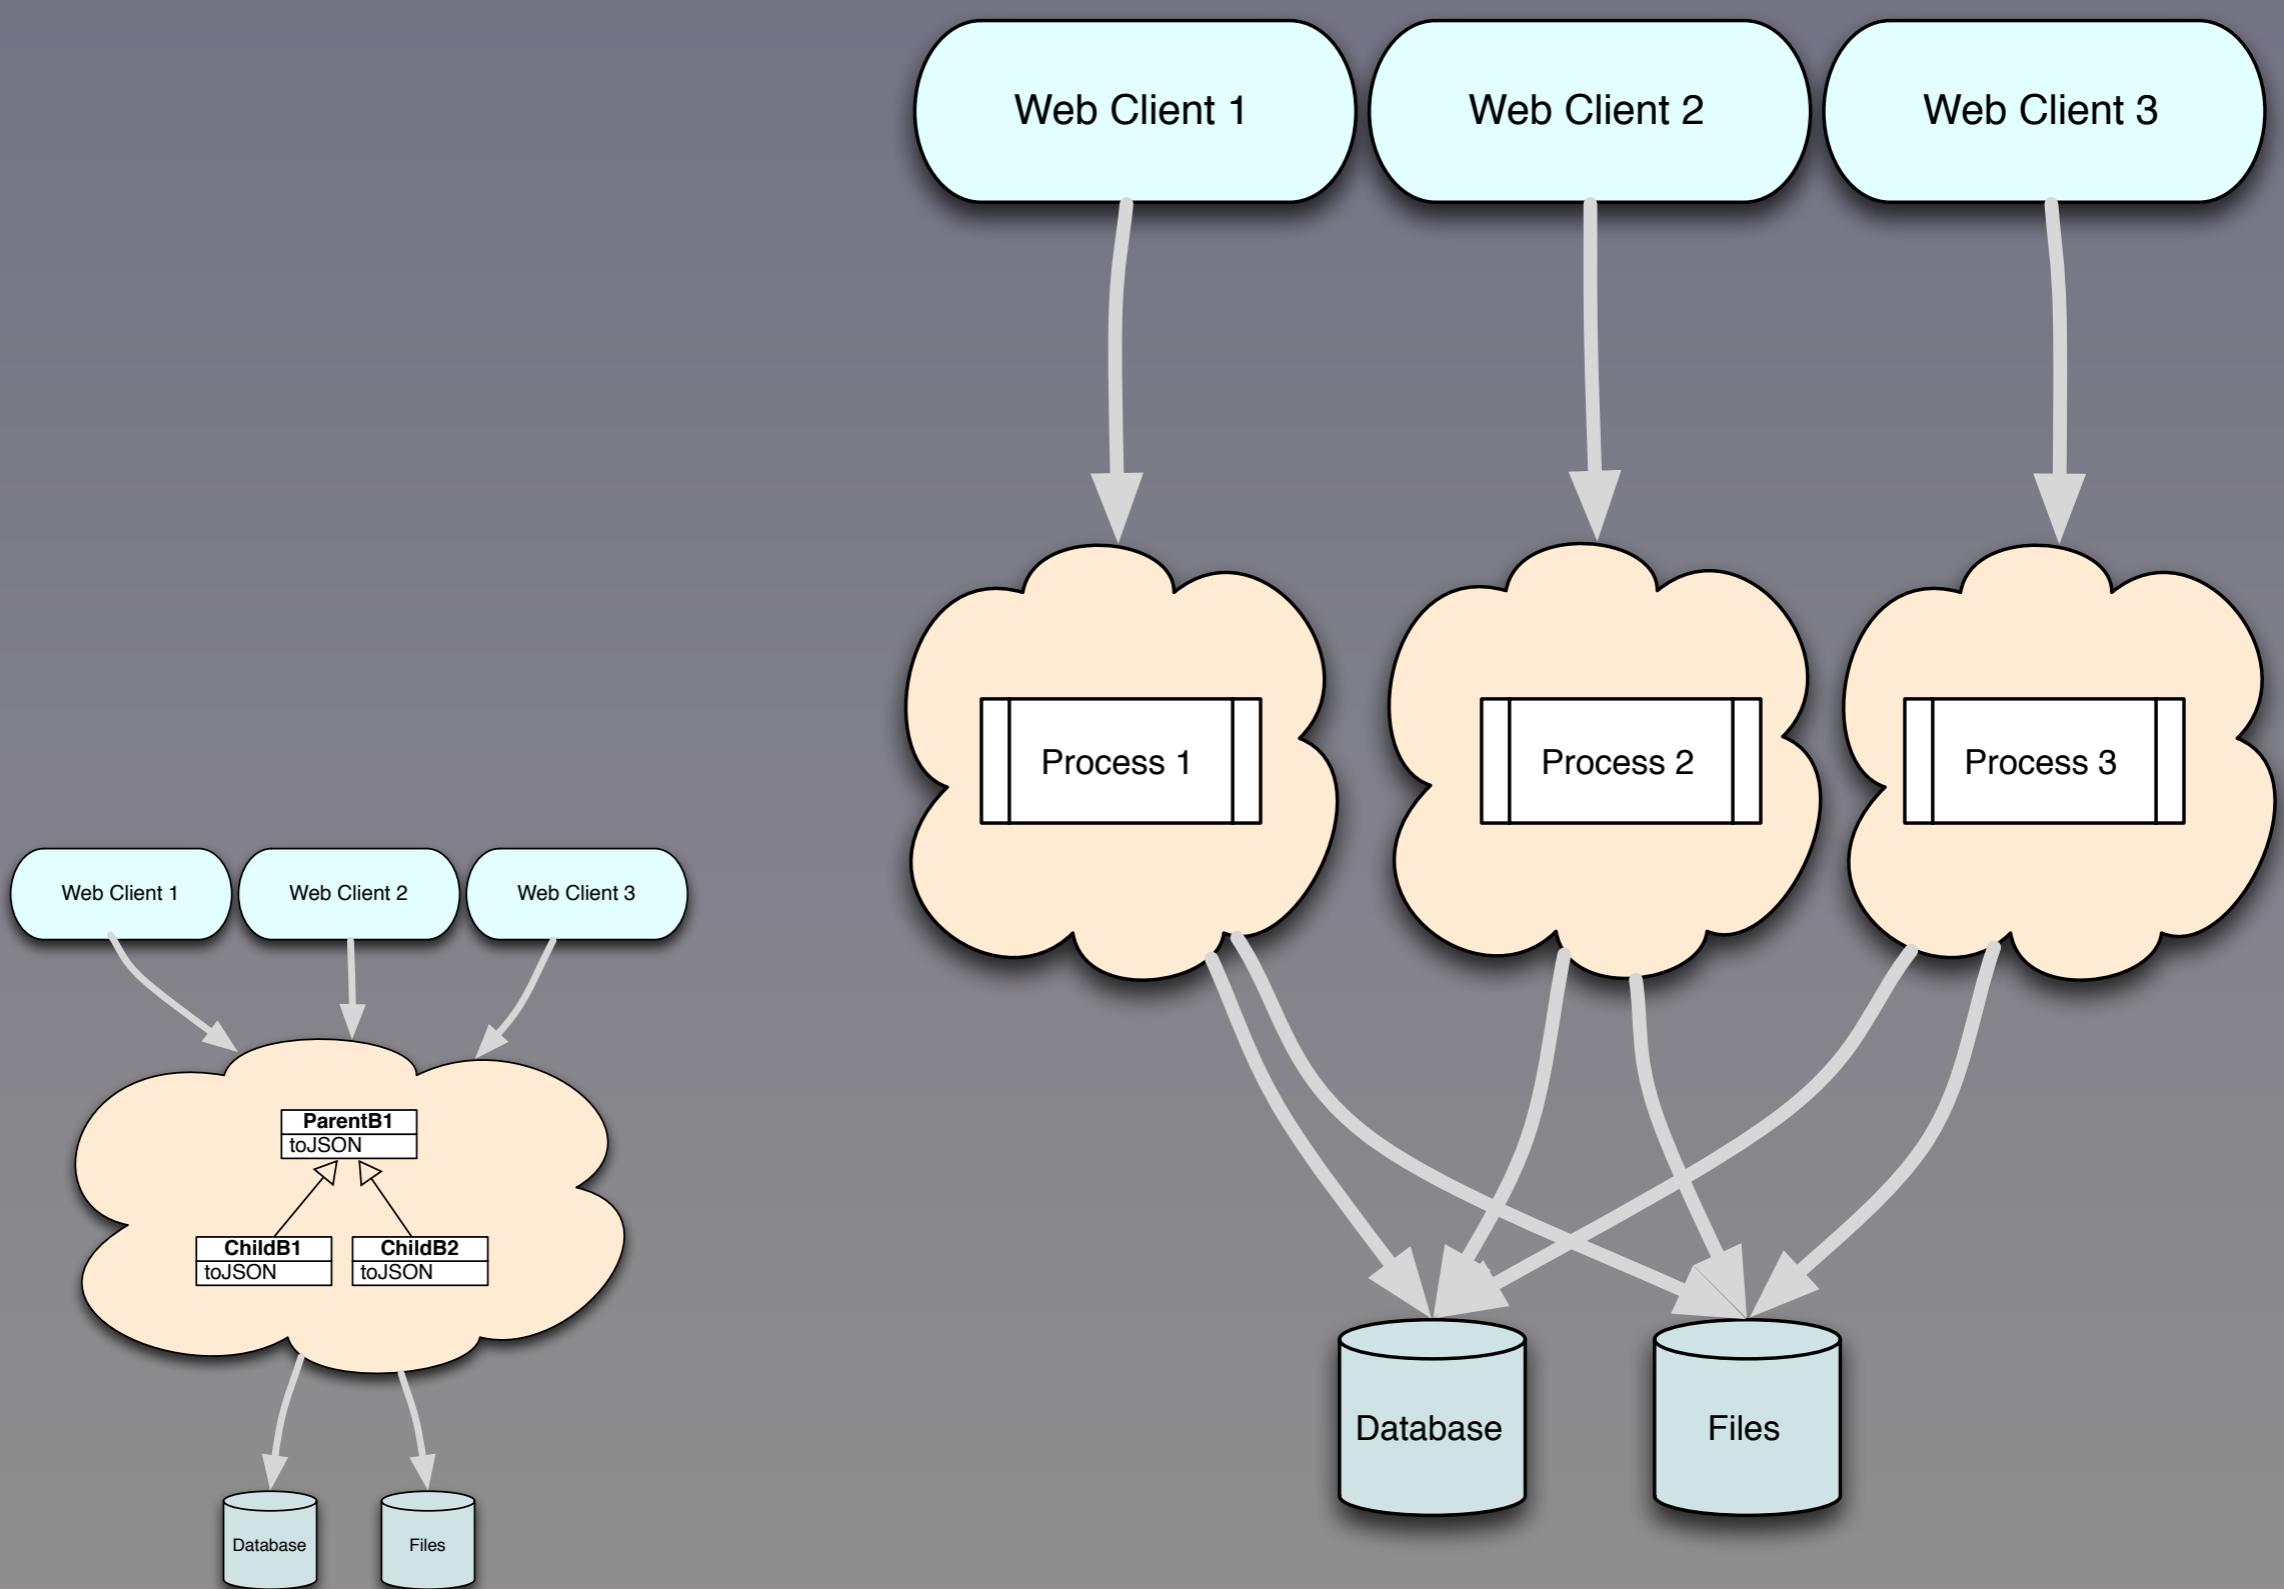
- Focus on:
  - Lists
  - Maps
  - Sets
  - Trees
  - ...



- NoSQL?
- Cassandra, HBase
- Riak, Redis
- MongoDB
- Neo4J
- • •





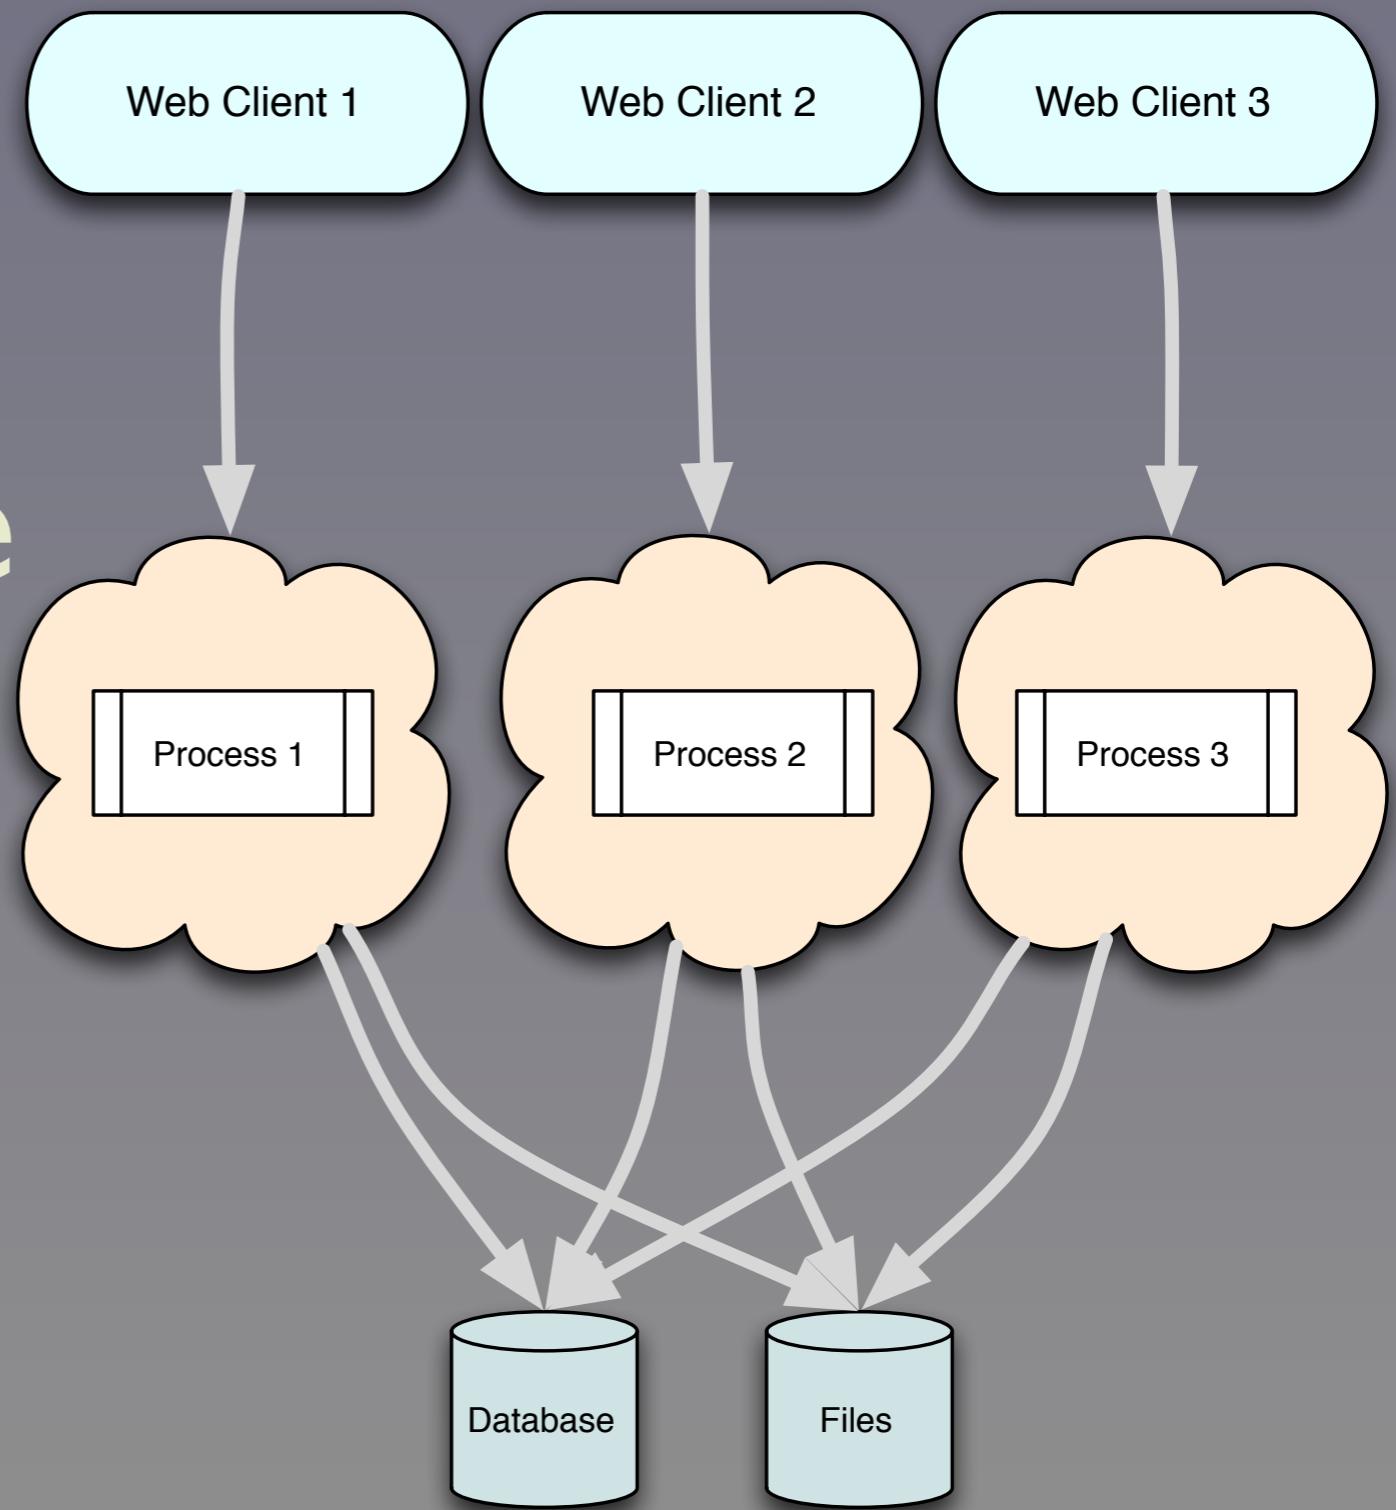


Sunday, June 9, 13

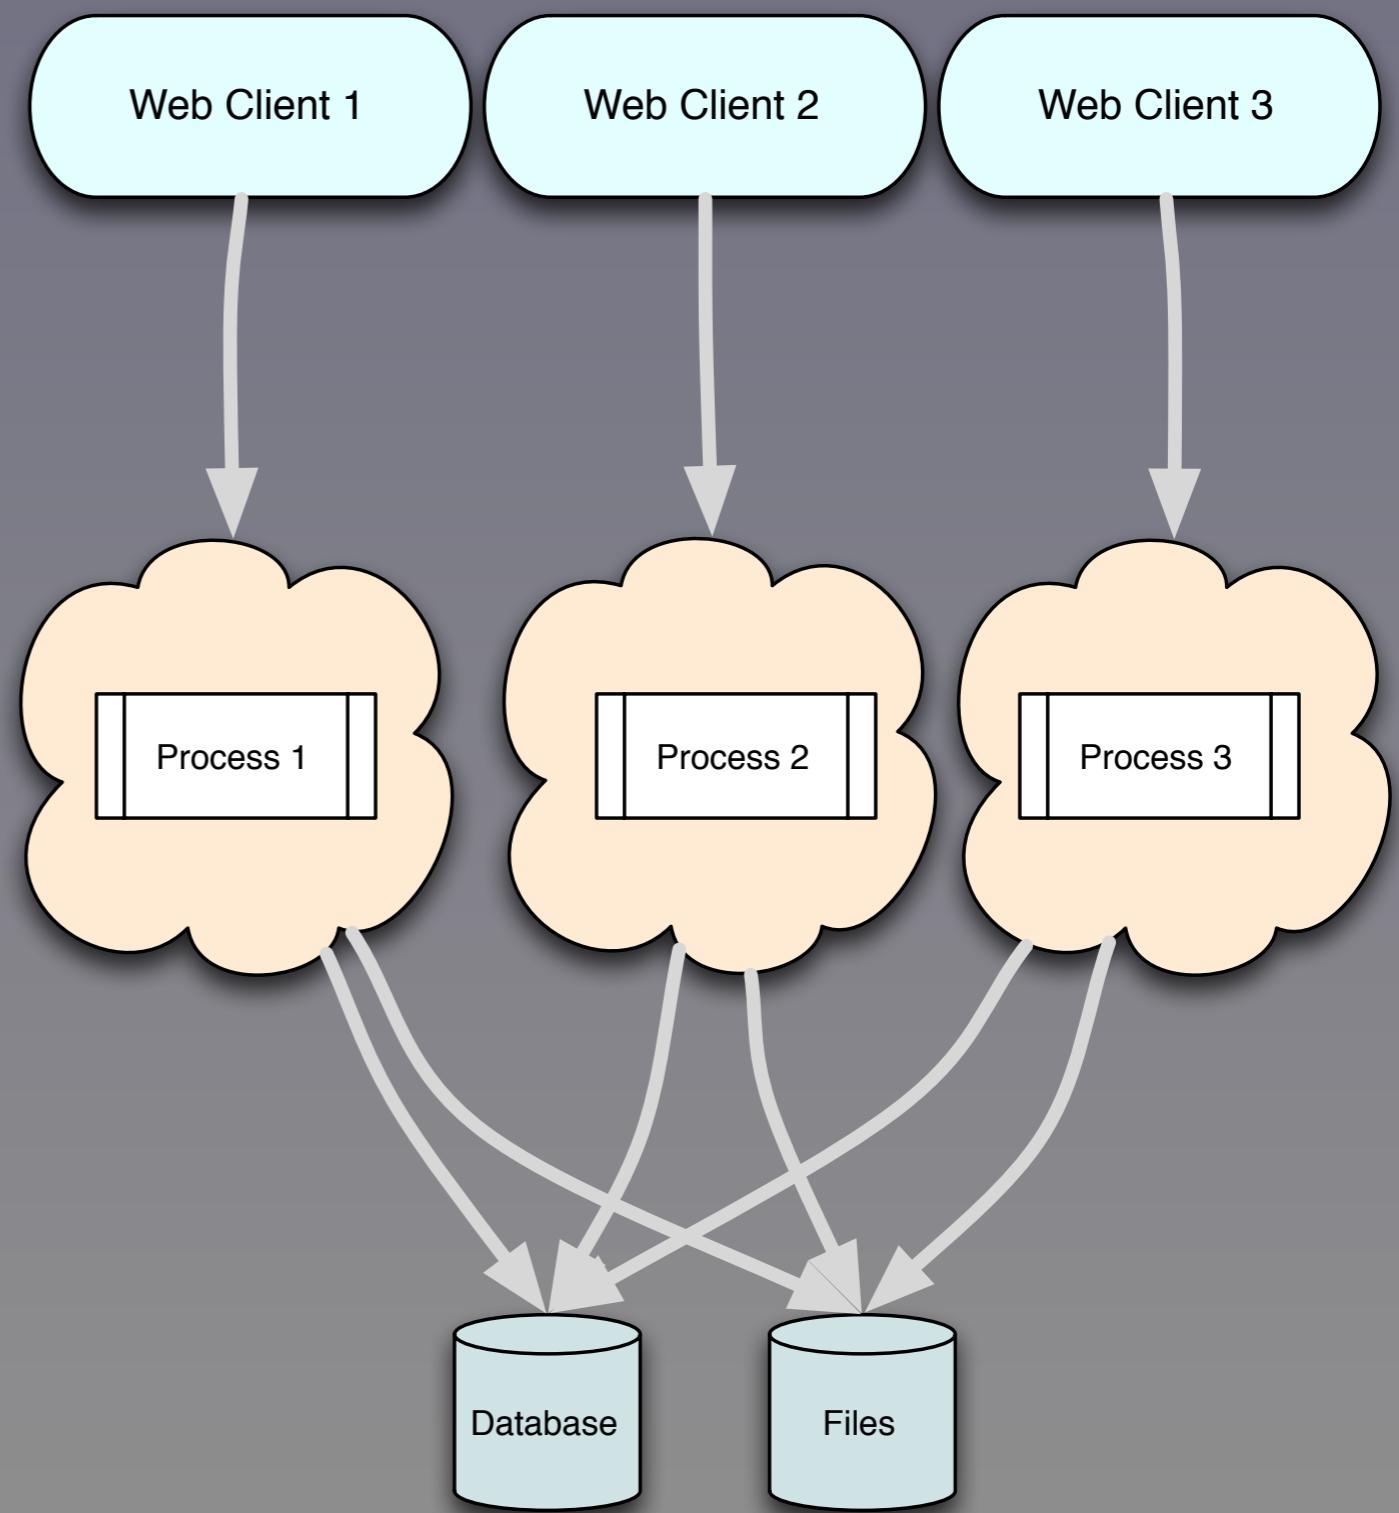
In a broader view, object models tend to push us towards centralized, complex systems that don't decompose well and stifle reuse and optimal deployment scenarios. FP code makes it easier to write smaller, focused services that we compose and deploy as appropriate. Each "ProcessN" could be a parallel copy of another process, for horizontal, "shared-nothing" scalability, or some of these processes could be other services...

Smaller, focused services scale better, especially horizontally. They also don't encapsulate more business logic than is required, and this (informal) architecture is also suitable for scaling ML and related algorithms.

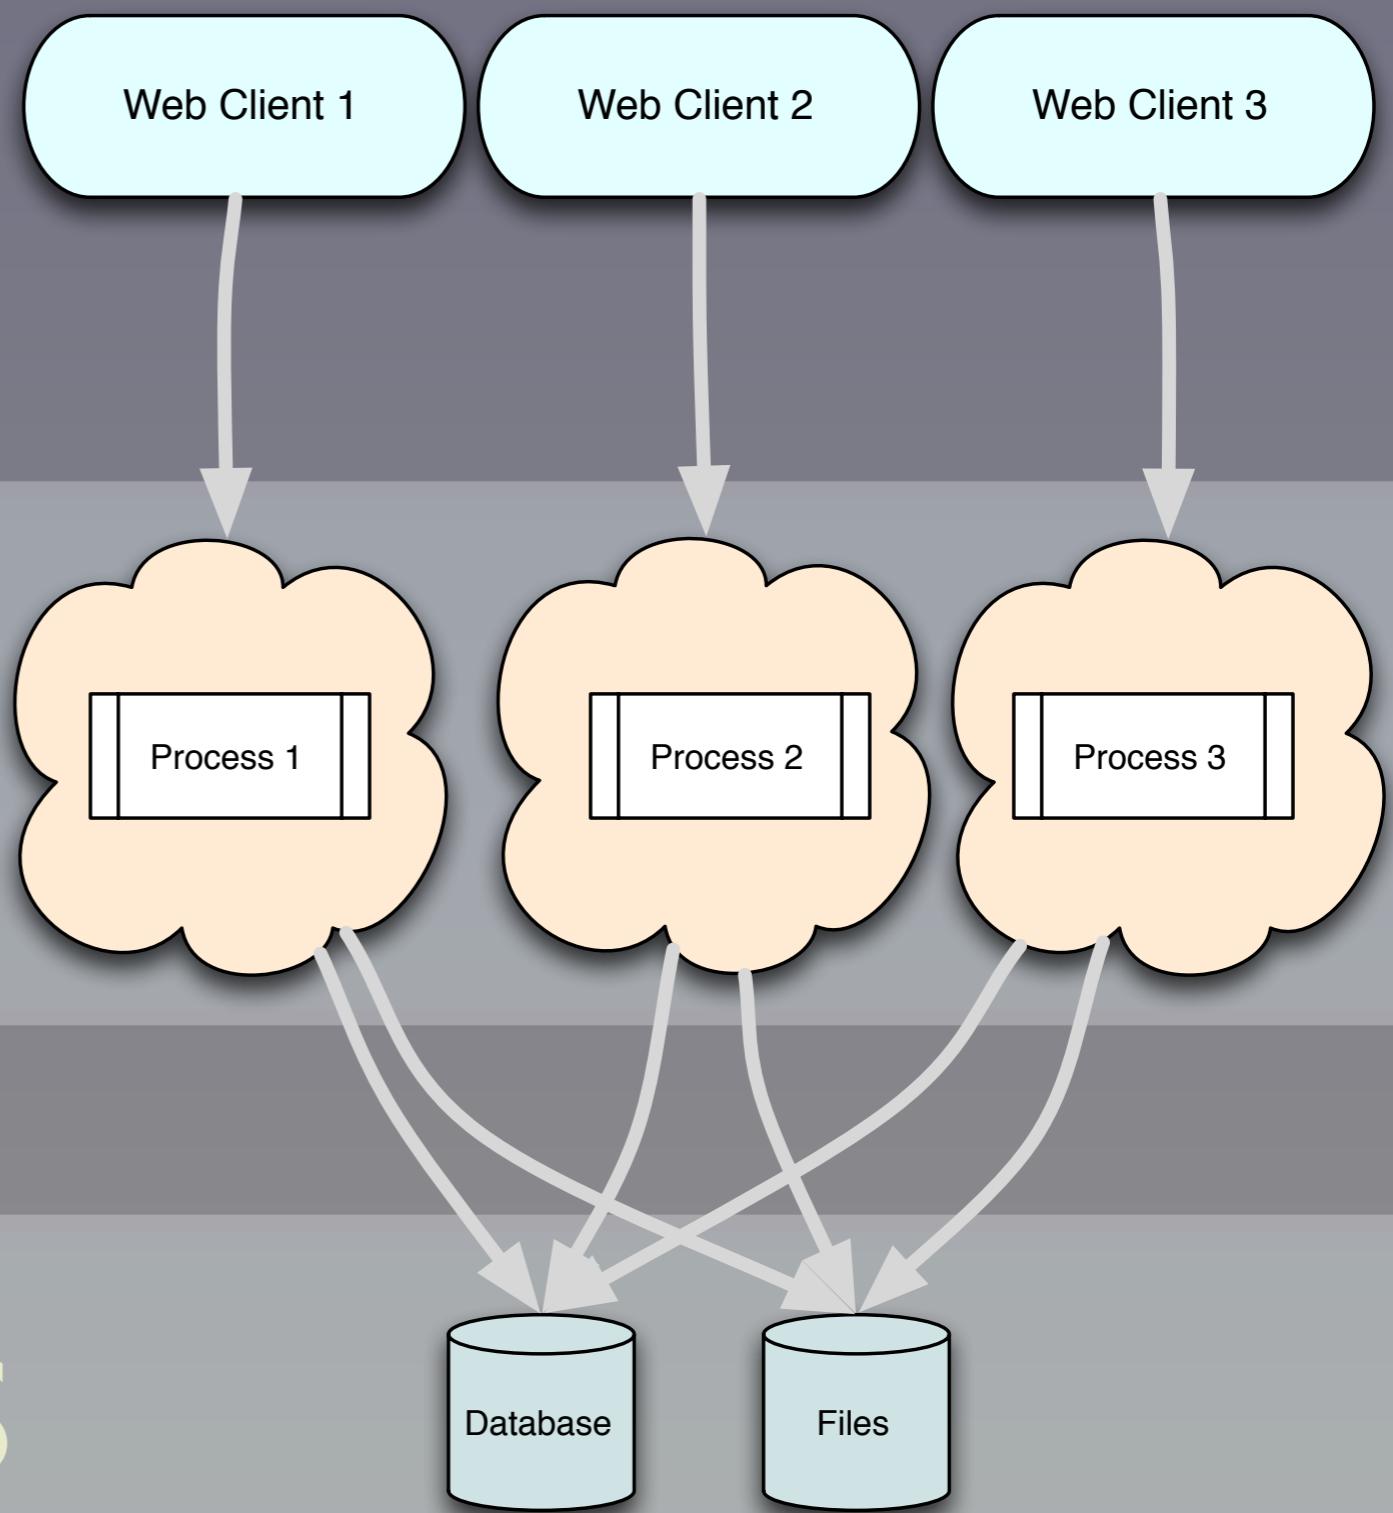
- Smaller, more focused middleware



- Data Size ↑
- Formal Schema ↓
- Data-Driven Programs ↑

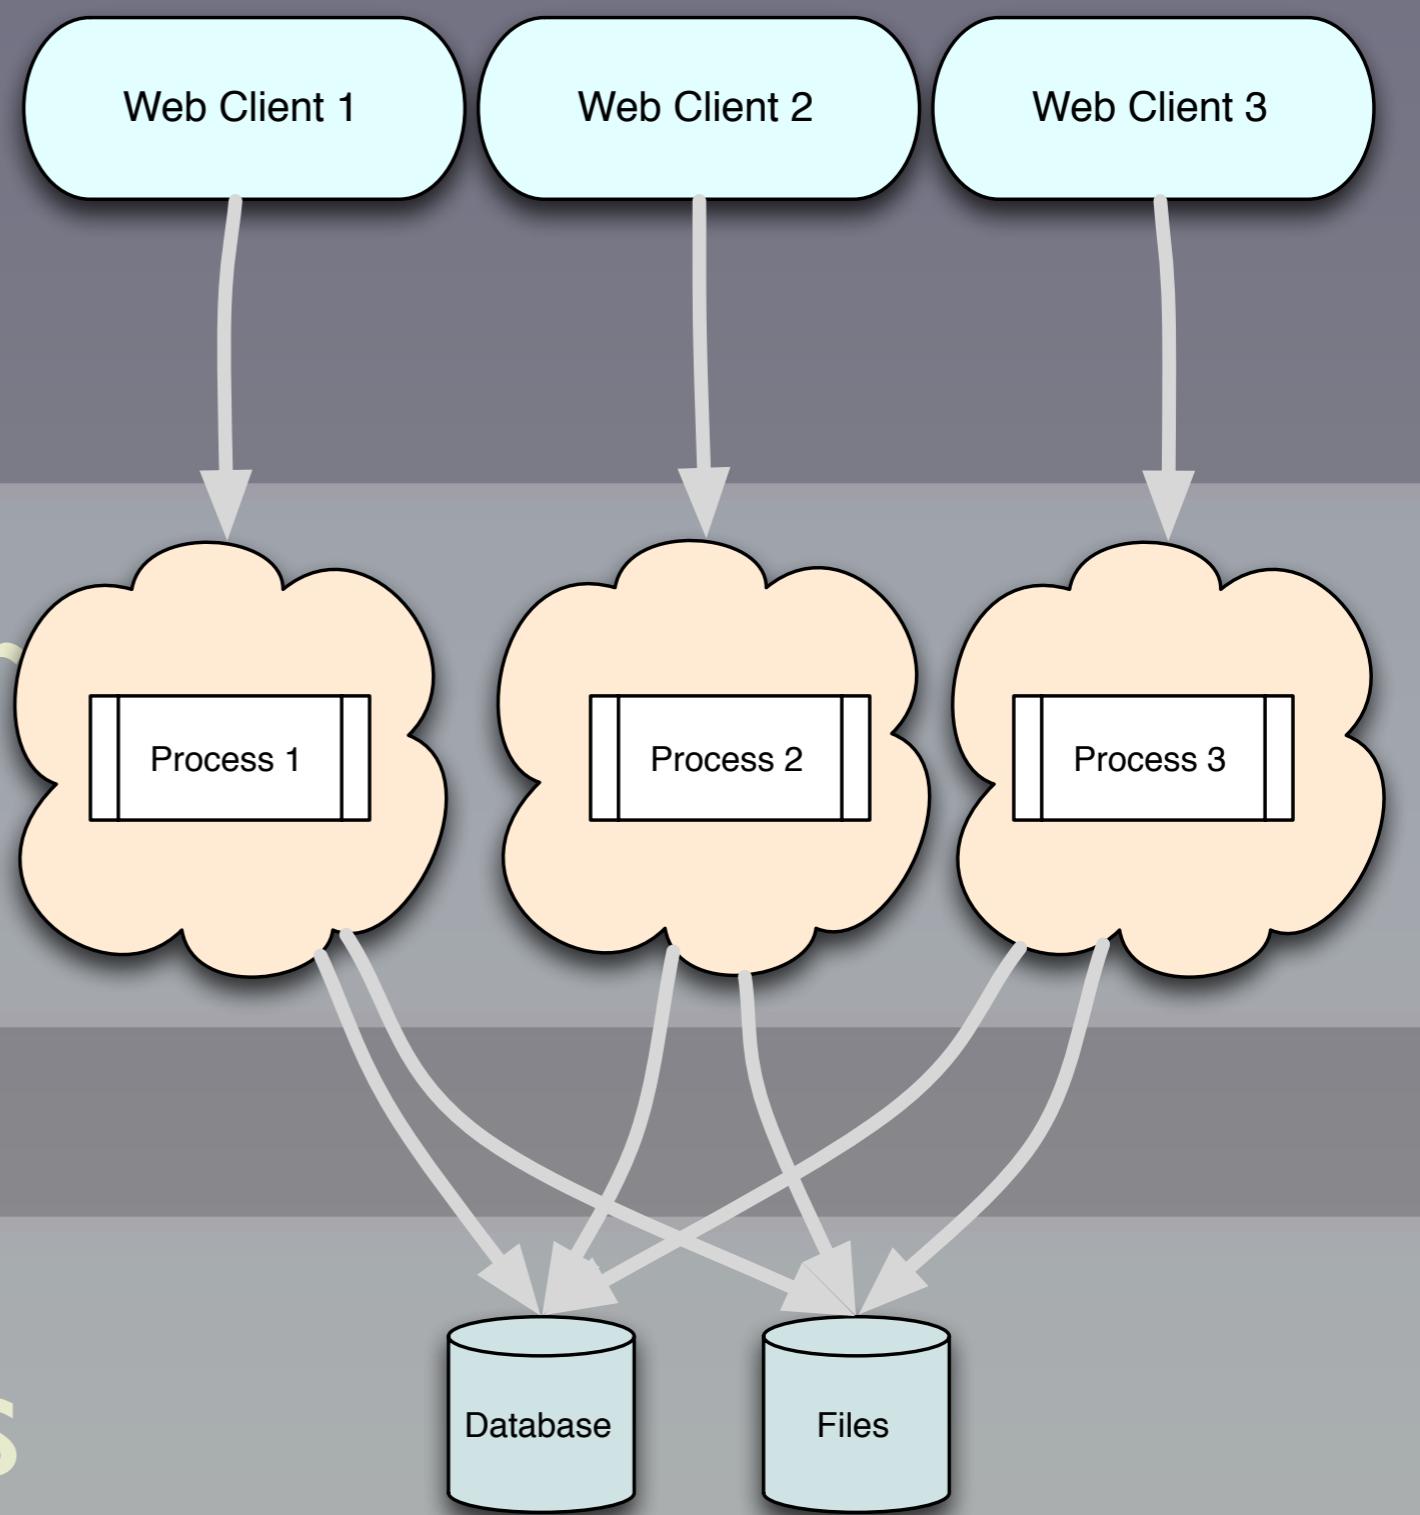


- MapReduce



- Distributed FS

- Hadoop, other middleware

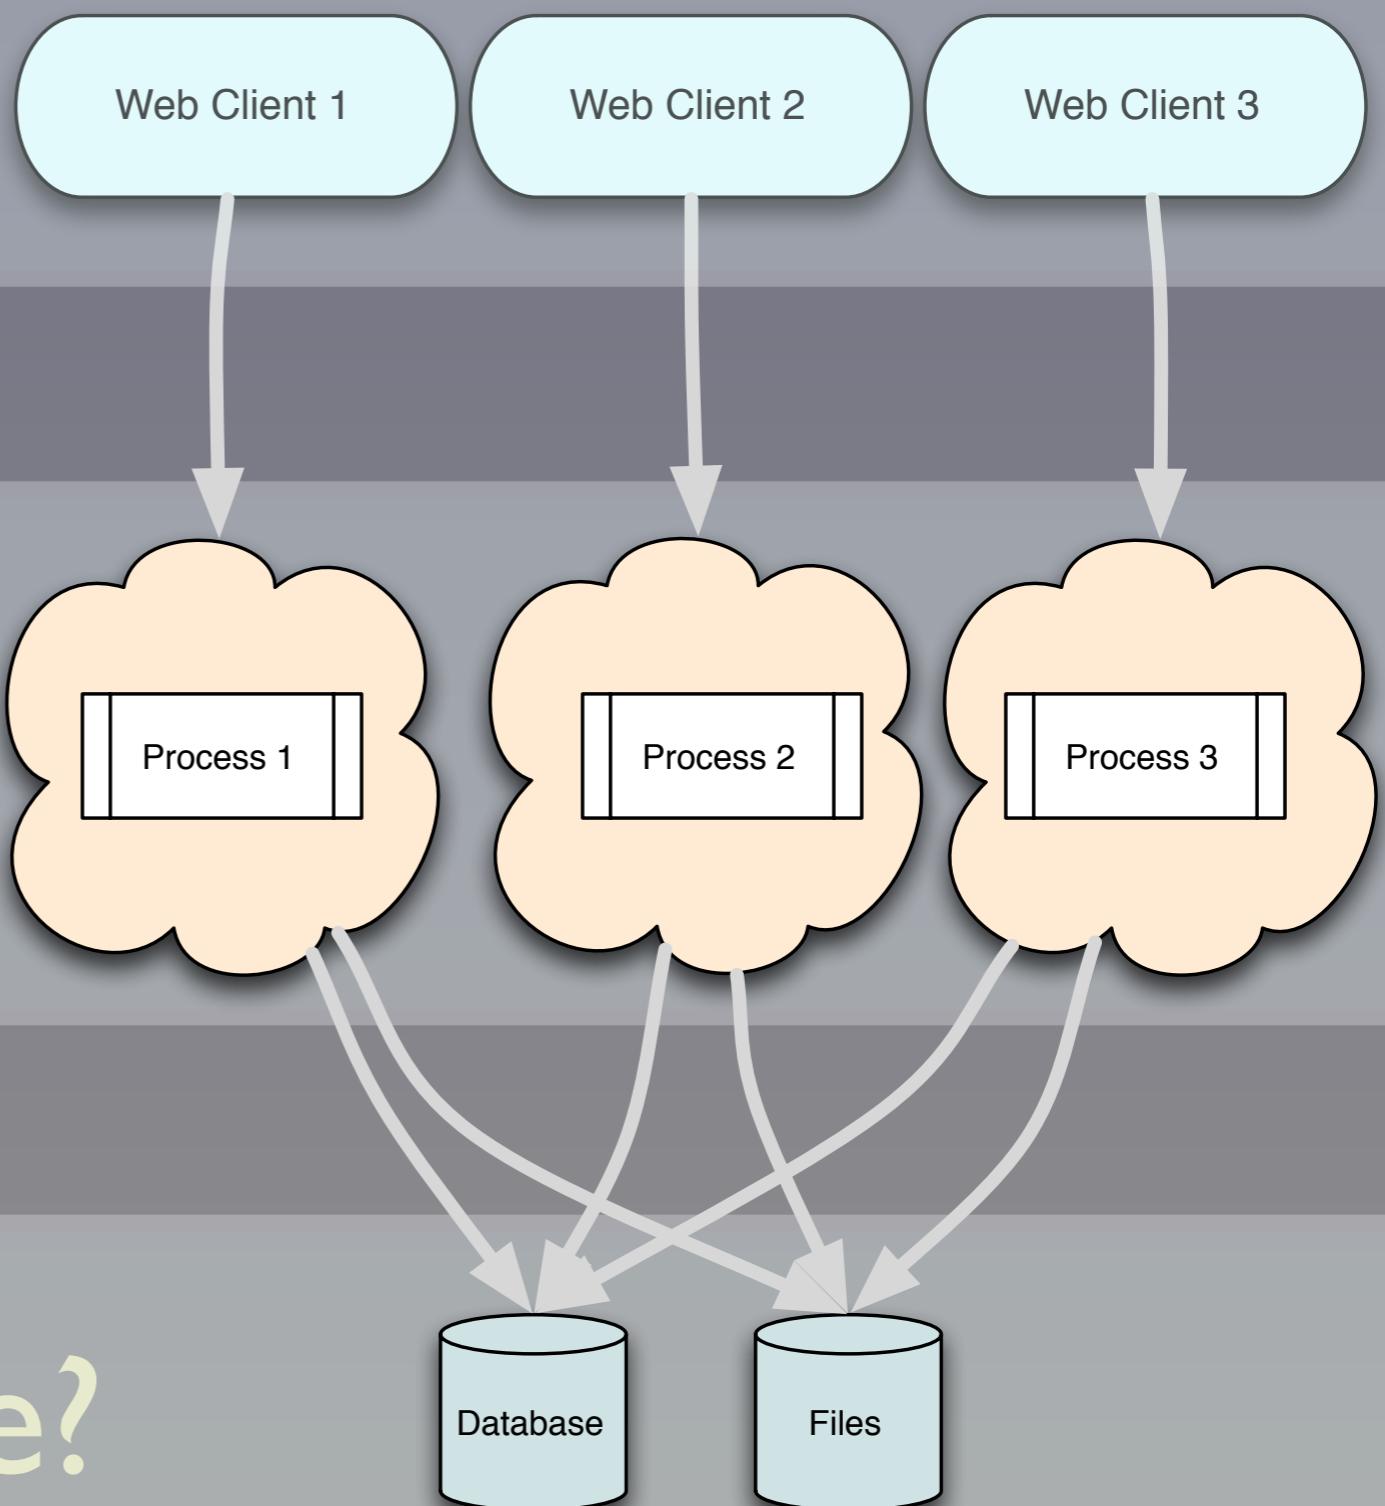


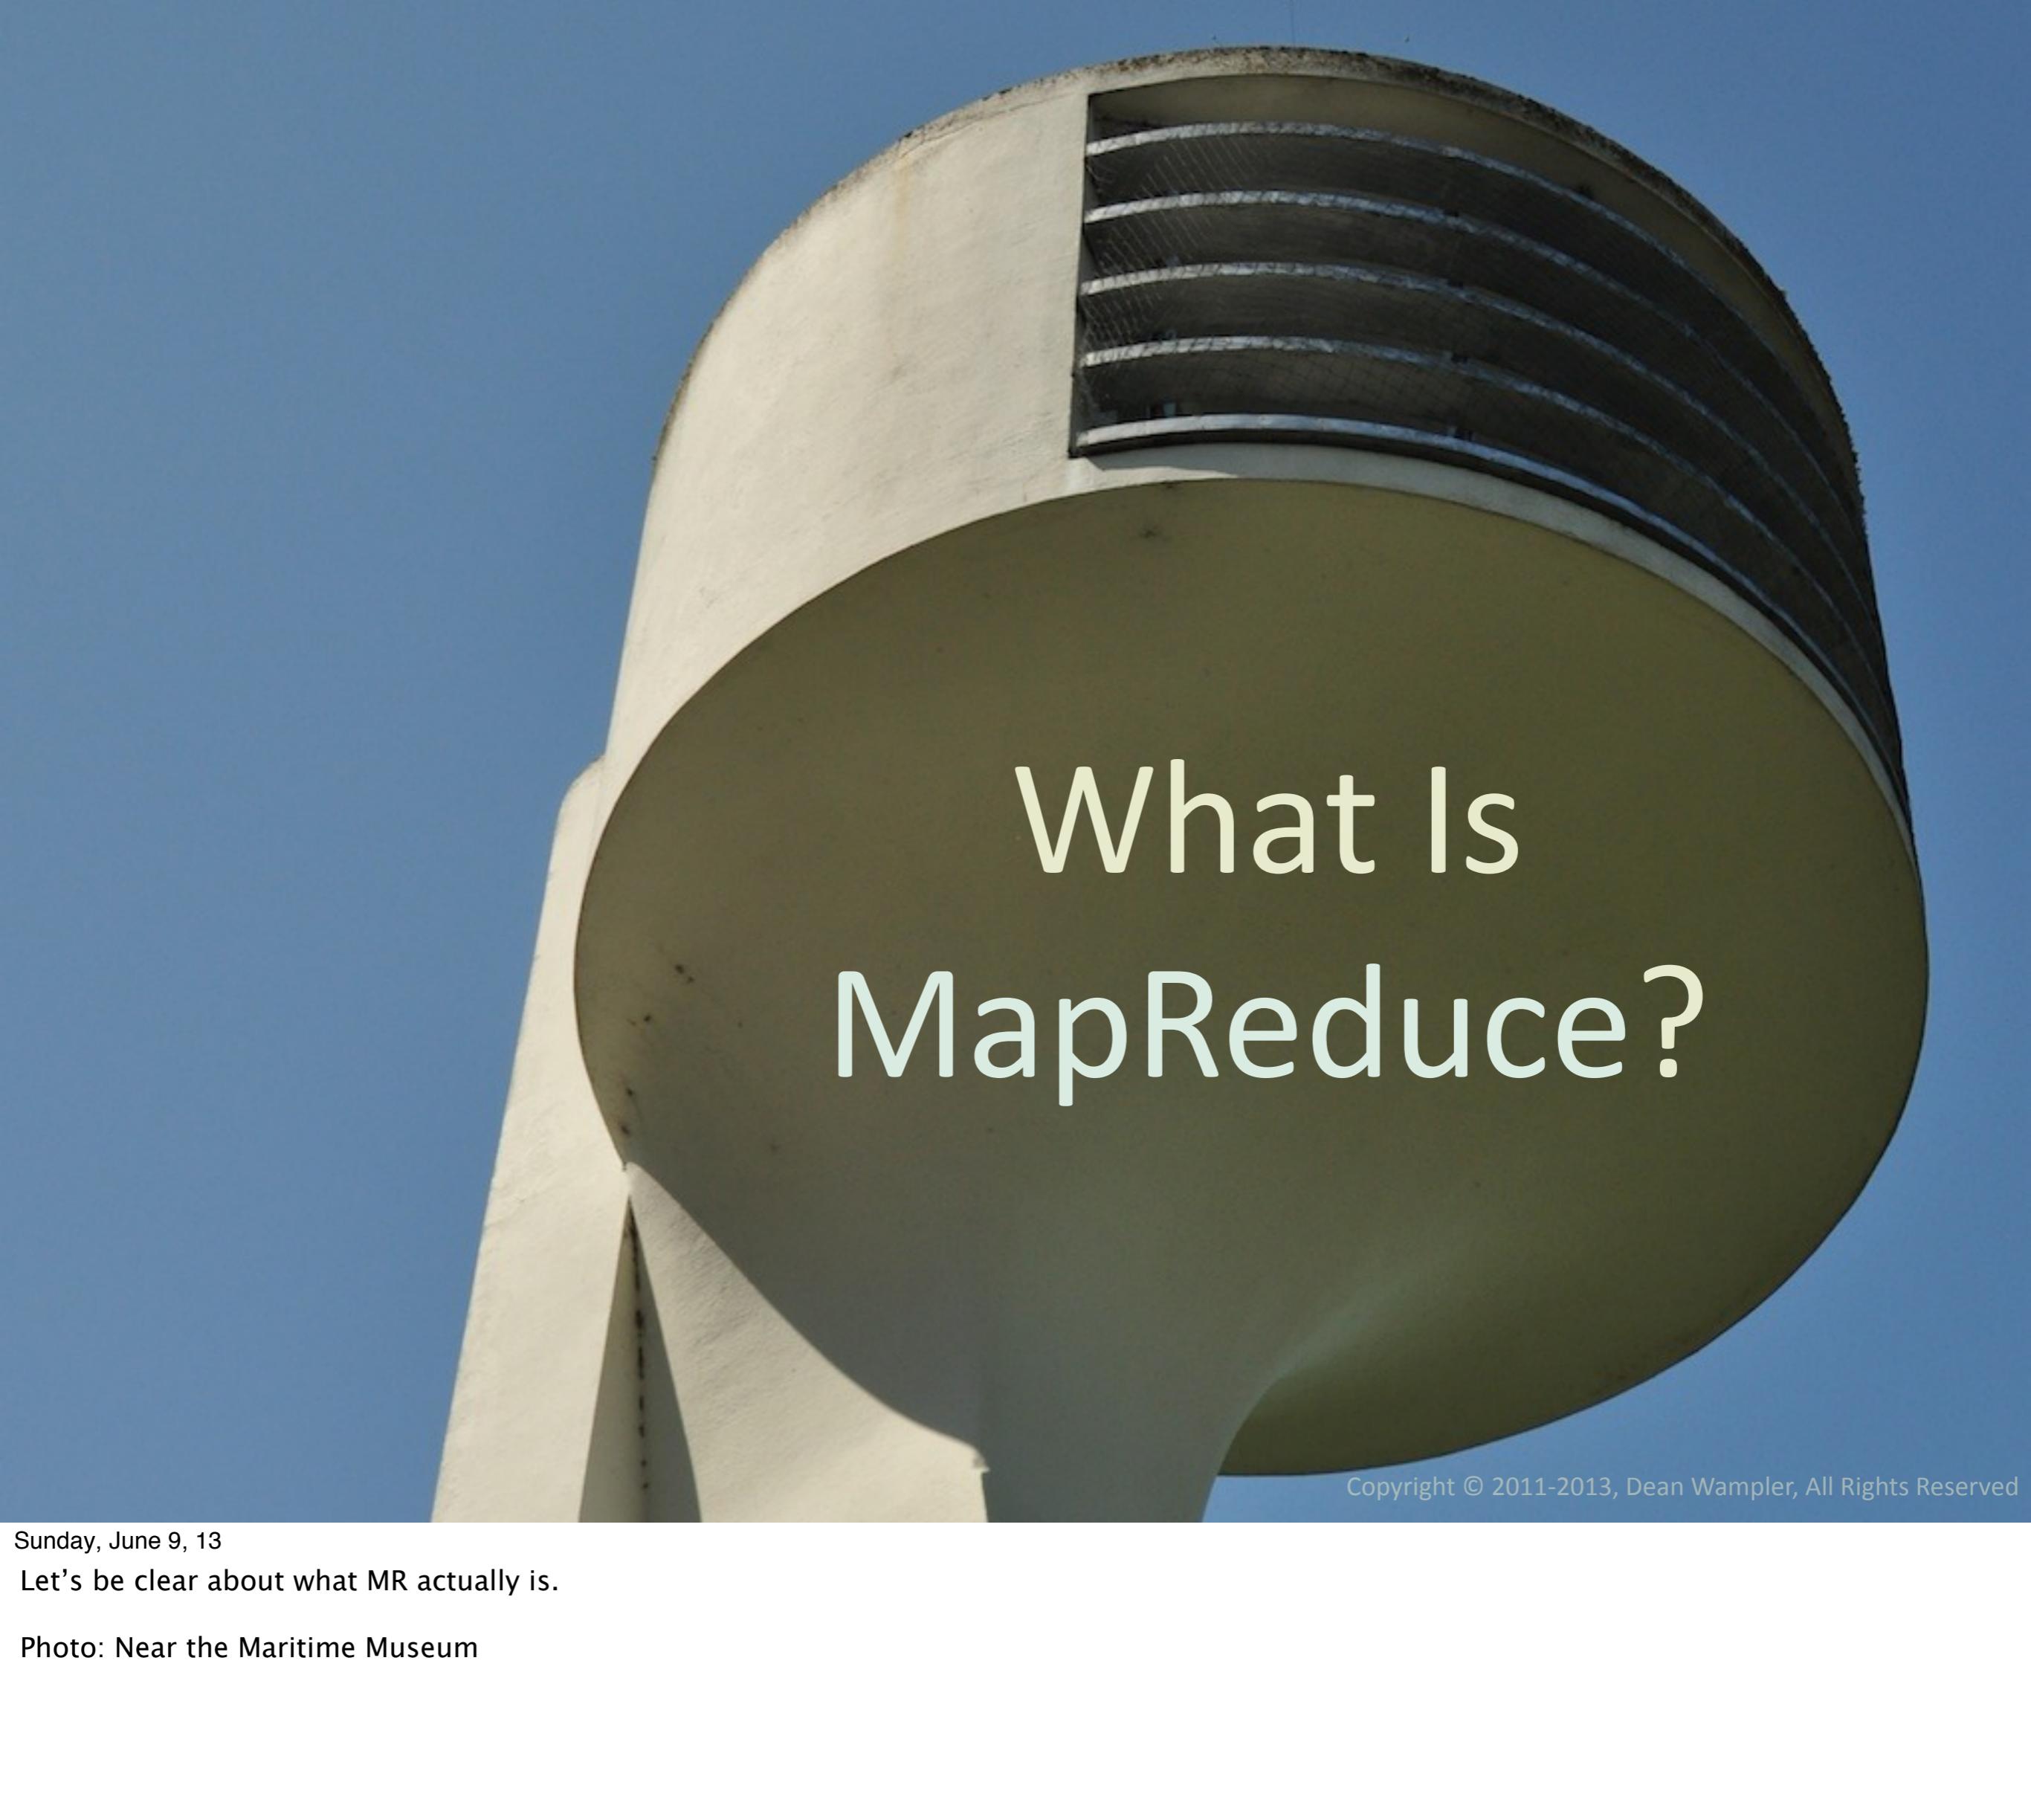
- NoSQL stores

- JSON

- Node.js?

- JSON database?





# What Is MapReduce?

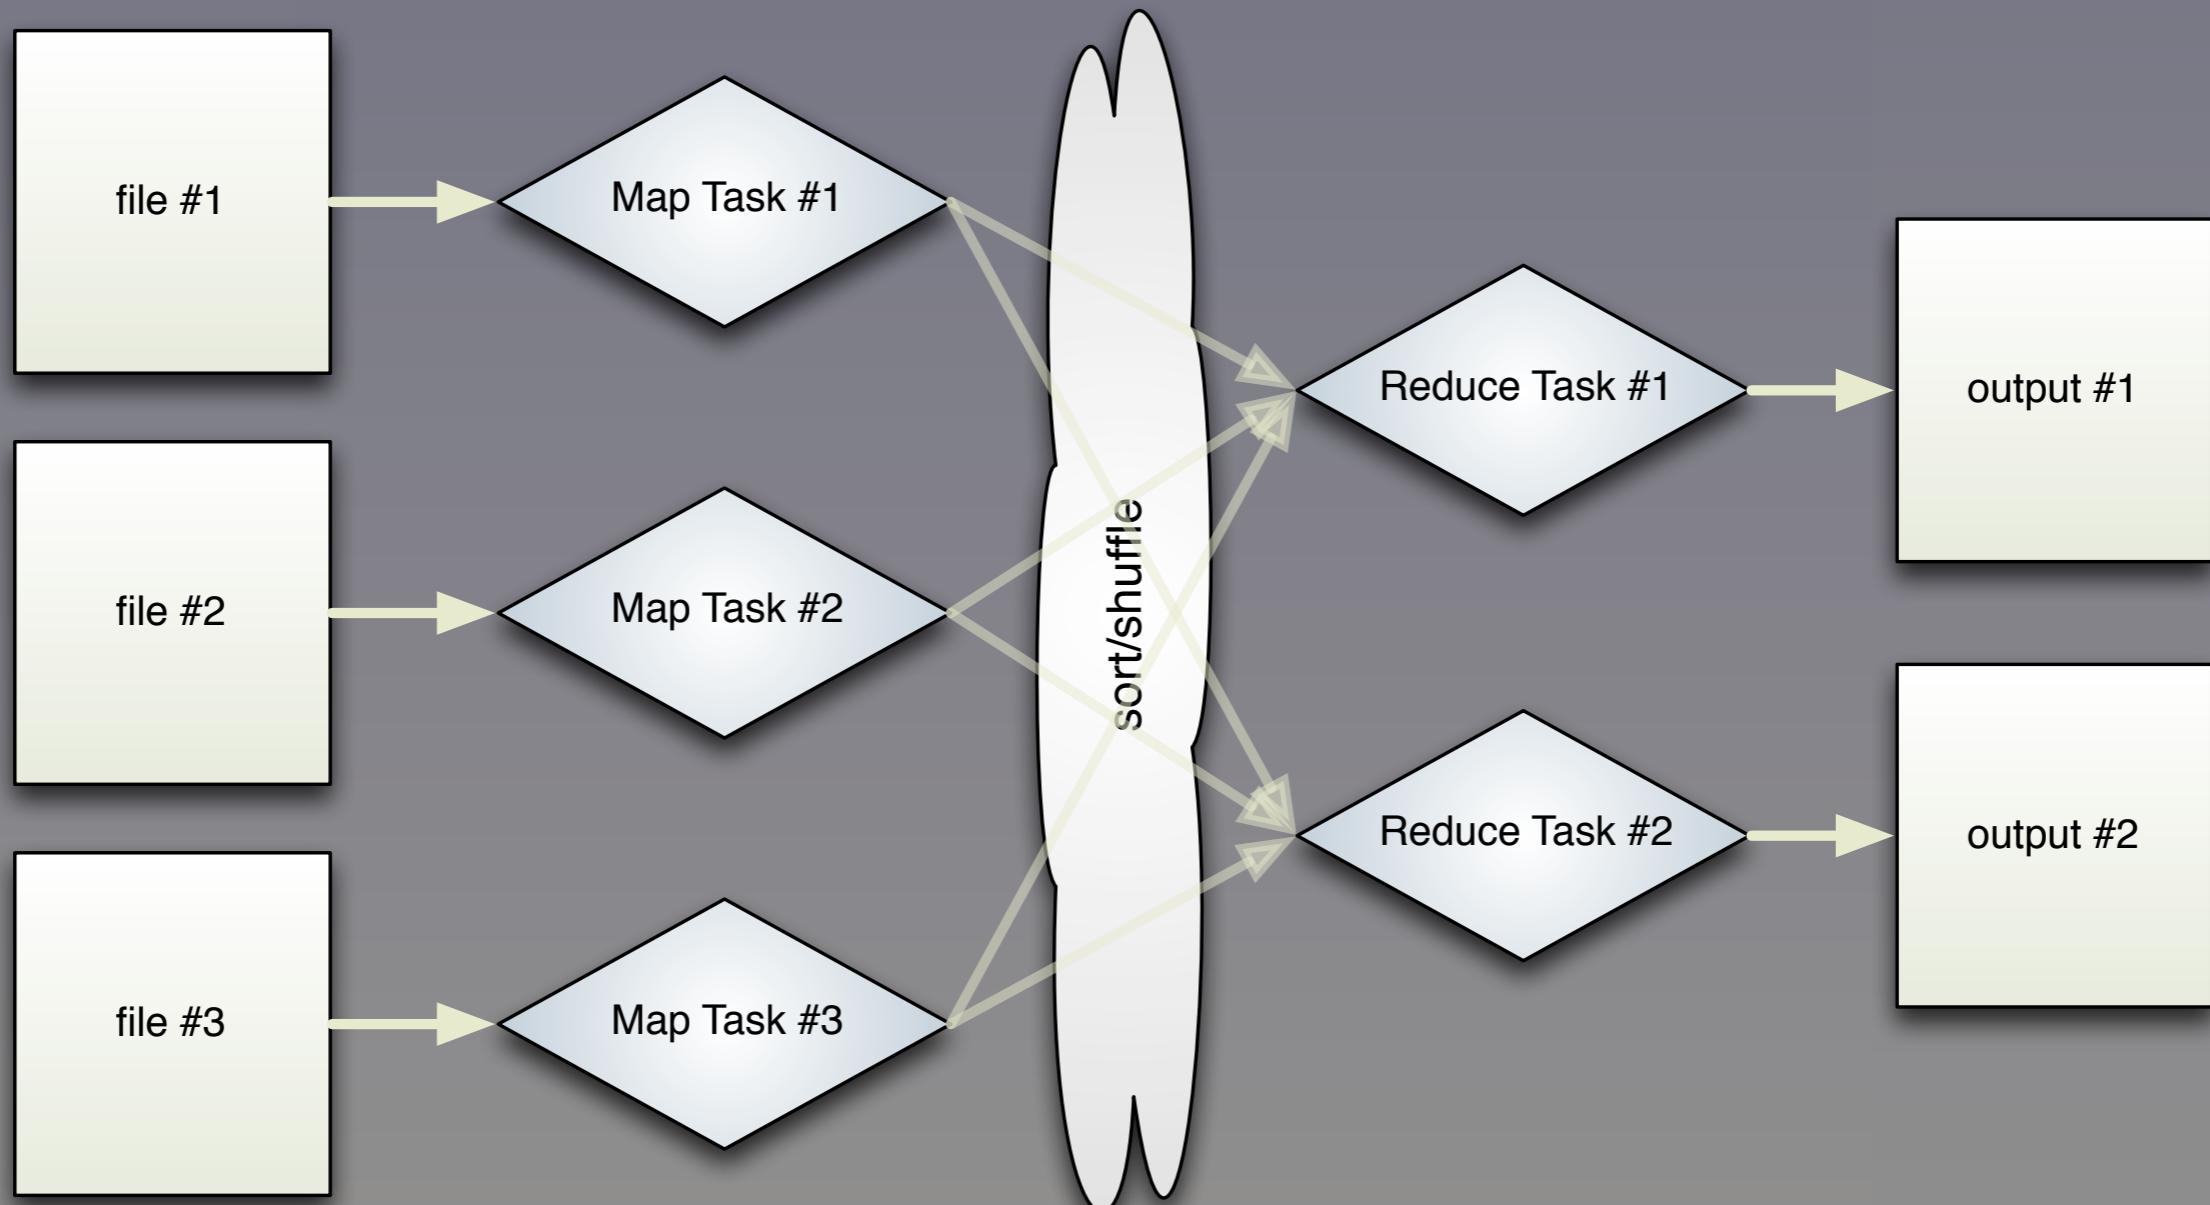
Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Let's be clear about what MR actually is.

Photo: Near the Maritime Museum

# A MapReduce Job



Each MR job reads in one or more (perhaps thousands...) of files, with one or more map tasks (JVM processes for each block of the files (where a block is  $n \times 64\text{MB}$ , for integer  $n$ , configured for HDFS)). The map task, which we write, tokenizes the input as required and outputs new key-value pairs, which are sorted inside each mapper, then shuffled to reducers, with all k-v pairs of the same key  $k$  go to the same reducer. A final reduction step is performed, then the output is dumped.

Arguably, Hadoop is our  
best, generic\* tool for  
scaling Big Data  
horizontally  
(at least today).

By design, Hadoop is  
*great* for *batch mode*  
data crunching.

*Not so great* for *event-stream* processing.

By design, Hadoop is  
*great for batch mode*  
data crunching.

*Not so great for  
transactions.*

MapReduce is very  
course-grained.

1-Map and  
1-Reduce phase...

*For Hadoop in  
particularly,  
the Java API is  
hard to use.*



# Use Cases

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Let's look at some use cases to understand the strengths and weaknesses of Hadoop-oriented solutions vs. NoSQL-oriented solutions.

Photo: Transamerica Building

# TL;DR

- *Hadoop*
- Very flexible compute model
- “Table” scans
- Batch mode
- NoSQL
- Focused on a particular data model
- Transactional
- Event driven

Top-level comparison. These are points of emphasis for these options. What I mean is that we tend to write Hadoop-centric or database-centric apps, depending on the use case.

# TL;DR

But *Hadoop MapReduce* is often used with a *NoSQL store*.

# Search



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Improving information search and retrieval, usually through some means of indexing. See the two search talks today by Ryan Tabora and Drew Raines.

# Lucene with Solr or ElasticSearch

*A specific solution  
for search.*

For very large data sets...

NetApp project:

Use *Solr* to store indices  
to data in *HDFS* and  
*HBase*.

# Low-cost Data Warehouse



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Data warehouse systems hit scalability limits and cost/TB concerns at large scale...

# Problem:

Your current *data warehouse* can only store *6-months* of data without a *\$1M upgrade*.

# Traditional DW

- Pros
- Mature
- Rich SQL, analytics
- Mid-size Data
- Cons
- Expensive - \$/TB
- Scalability limits

Data warehouses tend to be more scalable and a little less expensive than OLTP systems, which is why they are used to “warehouse” transactional data and perform analytics. However, their \$/TB is ~10x the cost on Hadoop and Hadoop scales to larger data sets.

# Solution?

Replace the  
*data warehouse*  
with *NoSQL*?

*SQL is very important  
for *data warehouse*  
applications.*

NoSQL does give you the more cost-effective storage, but SQL is very important for most DW applications, so your “NoSQL” store would need a powerful query tool to support common DW scenarios. However, DW experts usually won’t tolerate anything that isn’t SQL.

# Solution?

Replace the  
*data warehouse*  
with *Hadoop*?

- *Traditional DW*
  - + Mature
  - + Rich SQL, analytics
  - Scalability
  - \$\$/TB
- *Hadoop*
  - Less mature
  - + Improving SQL
  - + Scalable!
  - + Low \$\$/TB

Data warehouses tend to be more scalable and a little less expensive than OLTP systems, which is why they are used to “warehouse” transactional data and perform analytics. However, their \$\$/TB is ~10x the cost on Hadoop and Hadoop scales to larger data sets.

Hadoop has become a  
popular *data warehouse*  
supplement/replacement.

Many of my projects have offloaded an overburdened or expensive traditional data warehouse to Hadoop. Sometimes a wholesale replacement, but more often a supplemental strategy, at least for a transitional period of some duration.

# SQL on Hadoop

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Let's discuss the SQL options now, as I consider them very important. Otherwise, the majority of Data Analysts and casual SQL users would not find Hadoop very useful.

# Use SQL when you can!

- Hive: SQL on top of MapReduce.
- Shark: Hive ported to Spark.
- Impala: HiveQL with new, faster back end.
- Lingual: SQL on Cascading.

# Word Count in Hive SQL!

```
CREATE TABLE docs (line STRING);
LOAD DATA INPATH '/path/to/docs'
INTO TABLE docs;
```

```
CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
(SELECT explode(split(line, '\W+'))
 AS word FROM docs) w
GROUP BY word
ORDER BY word;
```

Works for Hive, Shark, and Impala

# Hive

- SQL dialect.
- Uses MapReduce back end.
  - So annoying latency.
- First SQL on Hadoop.
- Developed by Facebook.

# Shark

- HiveQL front end.
- Spark back end.
- Provides better performance.
- Developed by Berkeley AMP.

See <http://www.cloudera.com/content/cloudera/en/products/cloudera-enterprise-core/cloudera-enterprise-RTQ.html>. However, this was just announced a few ago (at the time of this writing), so it's not production ready quite yet...

# Impala

- HiveQL front end.
- C++ and Java back end.
- Provides up to 100x performance improvement!
- Developed by Cloudera.

# Lingual

- ANSI SQL front end.
- Cascading back end.
- Same strengths/weaknesses for runtime performance as Hive.

# Event Stream Processing



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Data warehouse systems hit scalability limits and cost/TB concerns at large scale...

Recall, Hadoop is *great*  
for *batch mode* data  
crunching.

*Not so great* for *event-  
stream* processing.



Storm!

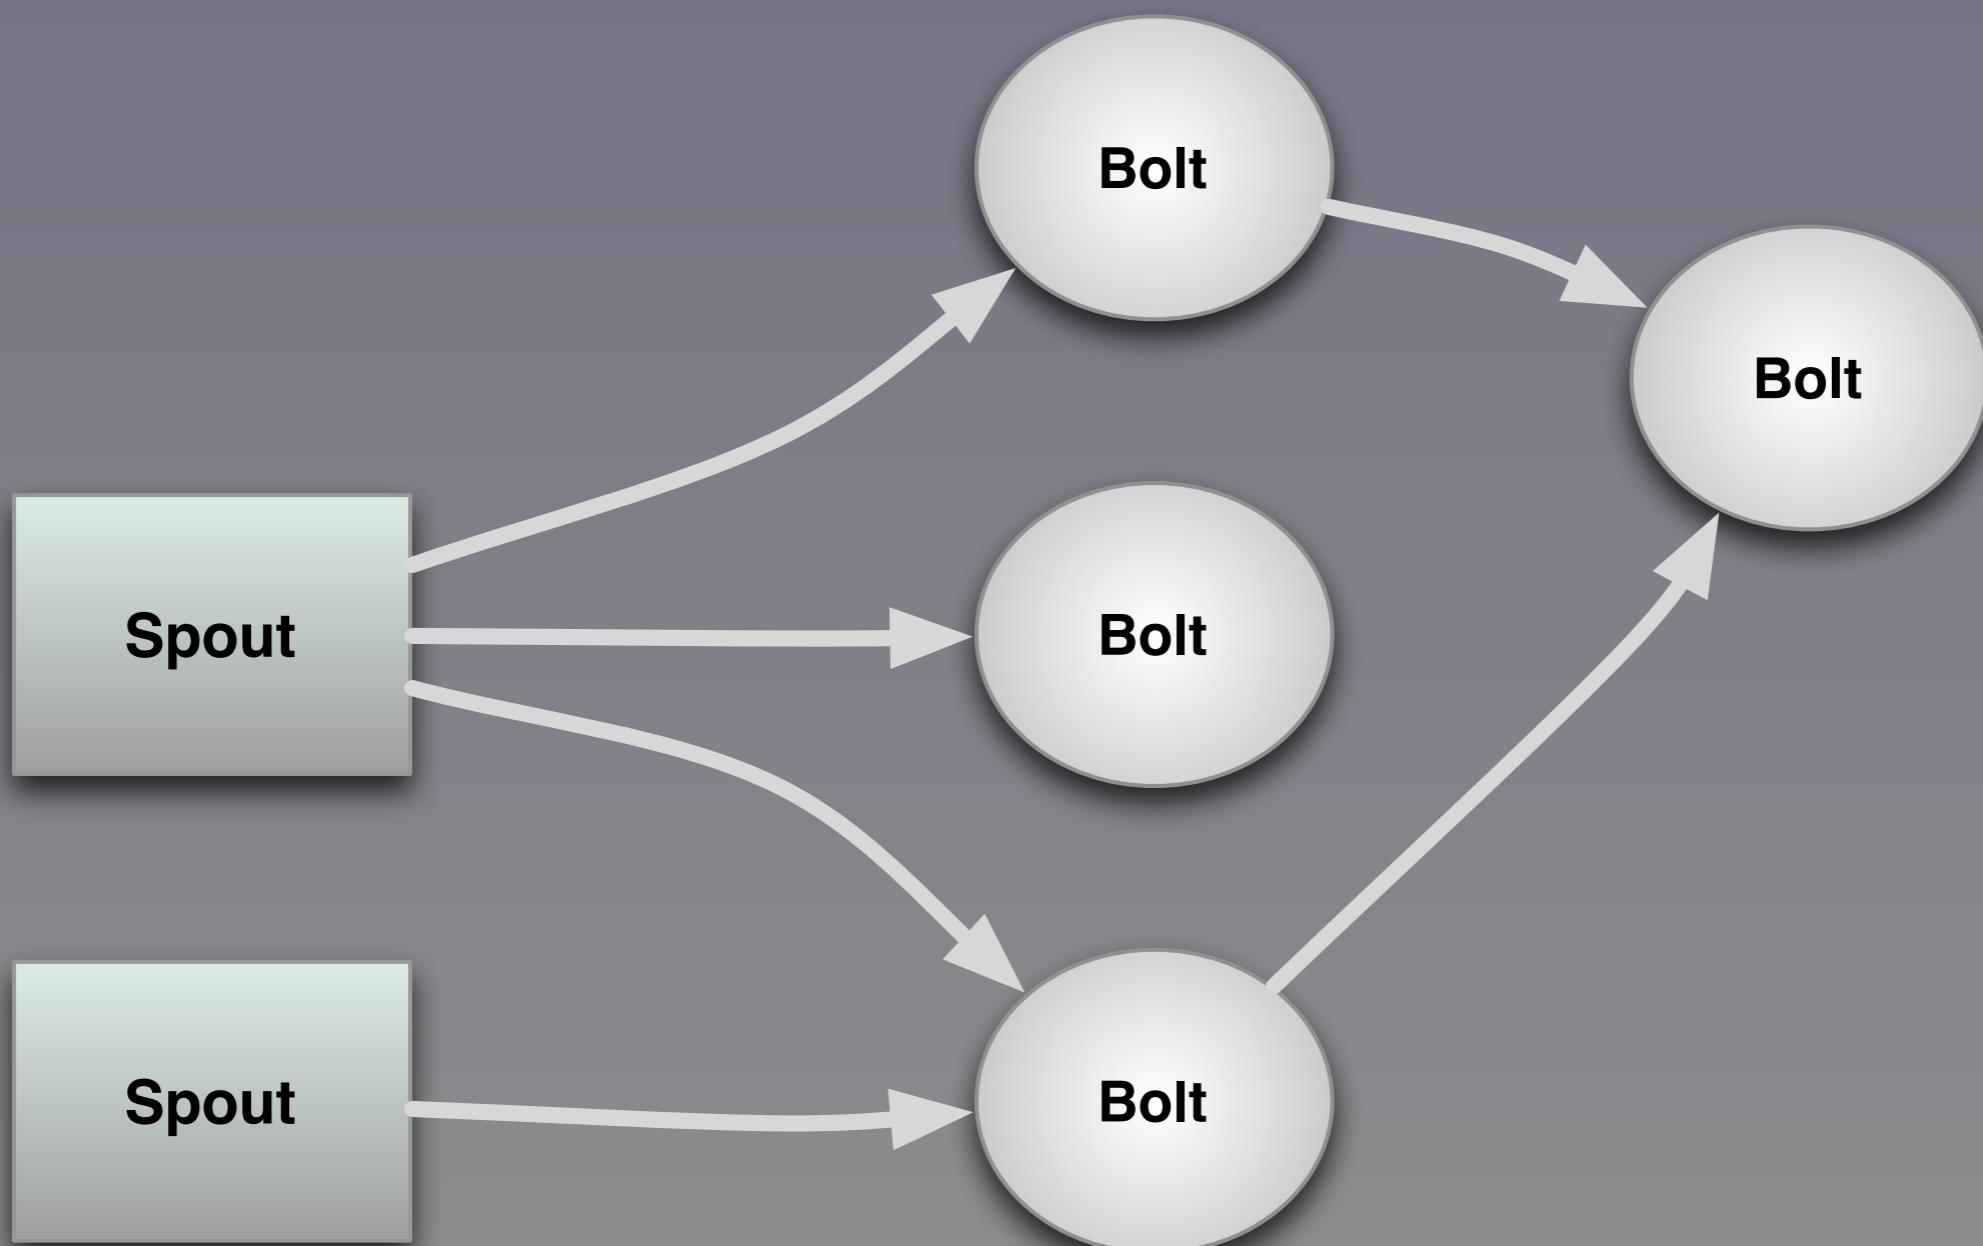
Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Nathan Marz's Storm system for scalable, distributed event stream processing. A complement to Hadoop, as he describes in detail in his book "Big Data" (Manning).

Photo: Top of the AON Building in Chicago after a Storm passed through.

Storm implements  
reliable, distributed  
event processing.



In Storm terminology, Spouts are data sources and bolts are the event processors. There are facilities to support reliable message handling, various sources encapsulated in Spouts and various targets of output. Distributed processing is baked in from the start.

- *Hadoop*
- + Cheap
- + Scalable
- + Commercial Support
- Batch mode
- *Storm*
- Less mature
- + Robust
- Commercial Support
- + “Real time”

If you need to respond to event streams, Hadoop simply isn't suitable. Storm is one possible solution. It was designed for this problem, but it is less mature and commercial support is just starting to appear.

# Databases to the Rescue?



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Photo: Actually, this is in Chicago, looking out my condo window.

# SQL or NoSQL Databases?

Databases are designed for fast, transactional updates.

So, consider a database for event processing.

# Machine Learning



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

ML includes recommendation engines (e.g., the way Netflix recommends movies to you or Amazon recommends products), classification (e.g., SPAM classifiers, character and image recognition), and clustering. Other specialized examples include text mining and other forms of natural language processing (NLP).

- *Recommendations*: Netflix movies, Amazon products, ...
- *Classification*: SPAM filters, character recognition, ...
- *Clustering*: Find groups in social networks, ...

Hadoop has a *general-purpose compute model*.

Arbitrary *ML algorithms*  
can be implemented  
using *MapReduce*.

Having a general-purpose computation framework (as opposed to storage model) is useful for implementing arbitrary algorithms, without having to worry about distribution and scalability boilerplate, like you might for an ad-hoc middleware layer.

However, recall that  
MapReduce is very  
course-grained...

# Pattern

A new toolkit for writing  
*models* in SAS, etc.,  
then run them on  
*Cascading.*

# Spark

An alternative to  
*MapReduce* with more  
*flexible* and *composable*  
abstractions.

# Spark

Originally designed for  
*Machine Learning.*

However you *train* the  
models...

*Train with Hadoop, etc.*  
*Serve requests with*  
*NoSQL store.*

A common model is to use Hadoop to train models (e.g., recommendation engine) over very large data set, then store the model in an event-processing system (NoSQL, Storm) to serve requests in real time. (Problem: how do you update the model to reflect new inputs? Rerun training periodically or use an “online” algorithm – out of scope here!)

# Social Network Data



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

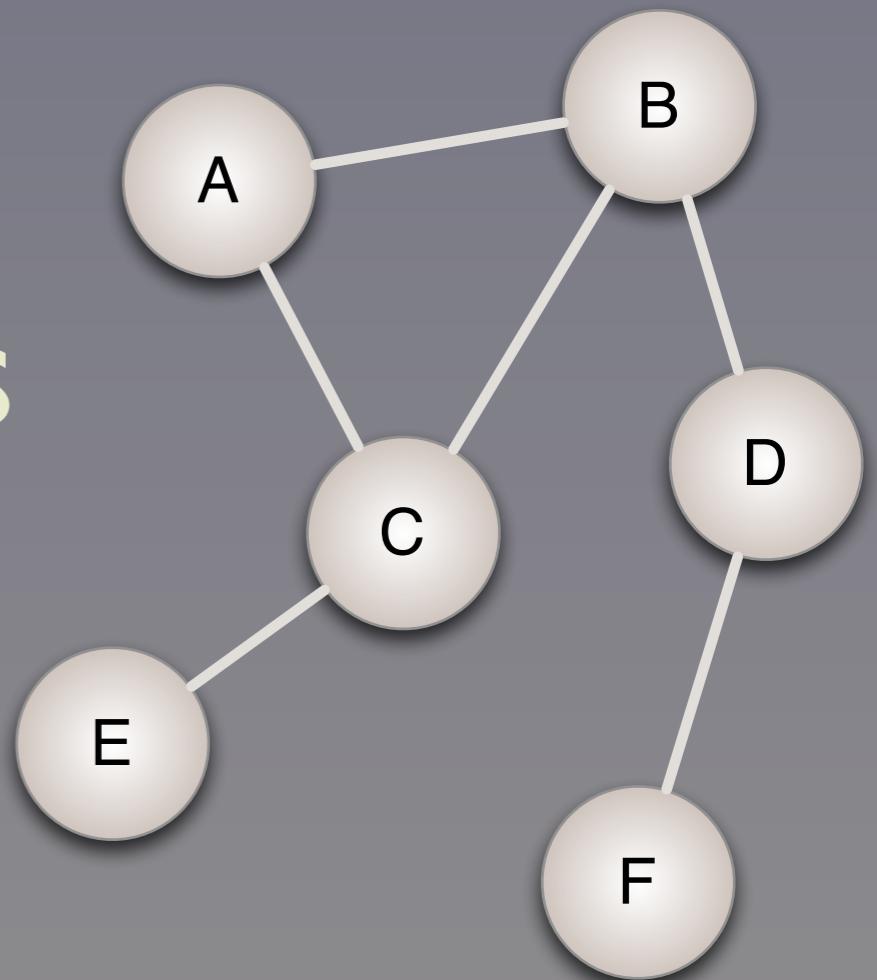
How should you represent associations in social networks, e.g., friends on Facebook, followers on Twitter, ...

# Google's Page Rank

Google invented MapReduce,  
... but MapReduce is not ideal for  
Page Rank and other graph  
algorithms.

# Why not MapReduce?

- 1 MR job for each iteration that updates all n nodes/edges.
- Graph saved to disk after each iteration.
- ...



# Use Graph Processing



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

A good summary presentation: <http://www.slideshare.net/shatteredNirvana/pregel-a-system-for-largescale-graph-processing>

Photo: San Francisco Bay

# Google's Pregel

- Pregel: New graph framework for Page Rank.
- Bulk, Synchronous Parallel (BSP).
  - Graphs are first-class citizens.
  - Efficiently processes updates...

Pregel is the name of the river that runs through the city of Königsberg, Prussia (now called Kaliningrad, Ukraine). 7 bridges crossed the river in the city (including to 5 to 2 islands between river branches). Leonhard Euler invented graph theory when we analyzed the question of whether or not you can cross all 7 bridges without retracing your steps (you can't).

# Open-source Alternatives

- Apache Giraph.
- Apache Hama.
- Aurelius Titan.

All are  
somewhat  
immature.

# Open-source Alternatives

- Neo4J.
- Mature, single machine graphical networks.

# Neo4J

- Not the same kind of distributed system like Pregel, but
  - More mature.
  - Commercially supported.
  - Maybe you don't need distributed?

# Batch Mode + Event-Stream Processing?

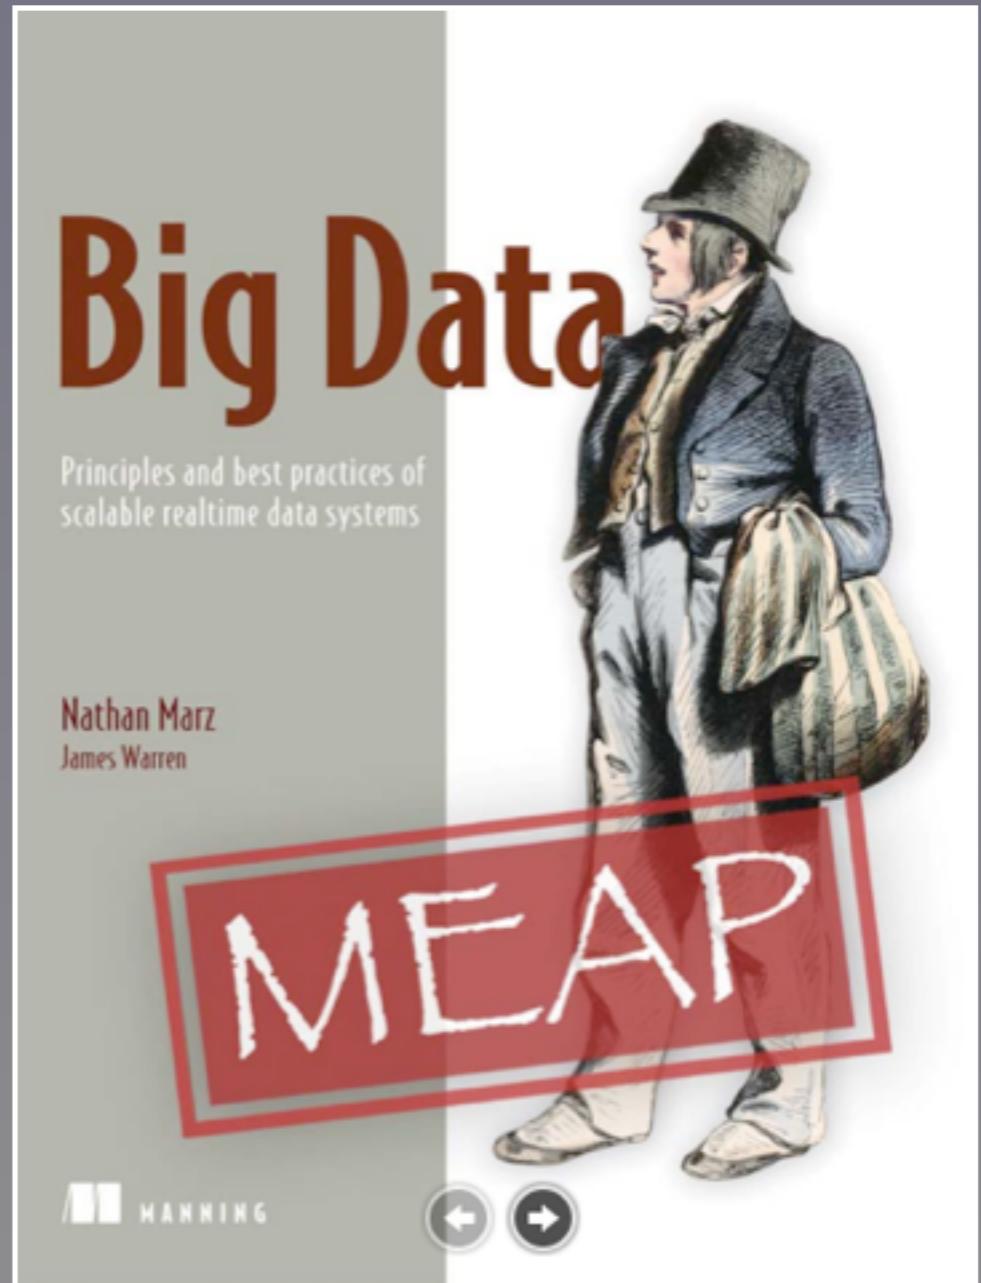
Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Now that we have cataloged some issues and solutions, let's recap and look forward.  
Photo: Grebe and distorted post reflection.

# “Big Data”

Nathan Marz’s  
vision for data  
systems...



- Use Hadoop for batch-mode processing of very large data sets.
- Use Storm for event handling, e.g., increment updates.
- Use NoSQL for flexible, scalable storage.
- Stir...

# So, where are we??



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Some thoughts about where the industry is and where it's headed.



Hadoop  
meets lots of  
needs now,  
but...

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13



# Hadoop MapReduce is the Enterprise Java Beans of our time.

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

I worked with EJBs a decade ago. The framework was completely invasive into your business logic. There were too many configuration options in XML files. The framework “paradigm” was a poor fit for most problems (like soft real time systems and most algorithms beyond Word Count). Internally, EJB implementations were inefficient and hard to optimize, because they relied on poorly considered object boundaries that muddled more natural boundaries. (I’ve argued in other presentations and my “FP for Java Devs” book that OOP is a poor modularity tool...) The fact is, Hadoop reminds me of EJBs in almost every way. It’s a 1st generation solution that mostly works okay and people do get work done with it, but just as the Spring Framework brought an essential rethinking to Enterprise Java, I think there is an essential rethink that needs to happen in Big Data, specifically around Hadoop. The functional programming community, is well positioned to create it...



I see Hadoop  
as  
*transitional.*

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

At least MapReduce. HDFS has longer legs, although it has its own flaws, like a single-point of failure NameNode that was only recently fixed (with replicated NNs), but is still immature, and it has upper bounds on storage. The MapR File System is superior in all these ways, but proprietary. Anyway, HDFS will last longer, but eventually be replaced by “next generation” distributed file systems, but MapReduce is already in decline and being replaced by more performant engines such as Impala, Spark, and Storm (for different purposes...).

# What about NoSQL??



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Does it feel “dated” like Hadoop?

The “good” NoSQL  
systems have all the  
right qualities, like  
minimal, focused  
abstractions.

# Purpose-built Tools



Sunday, June 9, 13

I see that a trend where completely generic tooling is giving way to more “purpose-built” tooling...

# ElasticSearch and Solr for Search.

New Hadoop  
file formats, optimize  
access (e.g., Parquet).

# New compute engines in Hadoop that replace MapReduce.

In the quest for ever better performance over massive datasets, the generic file formats in Hadoop and MapReduce are hitting a performance wall (although not everyone agrees). Parquet is column oriented & contains the data schema, like Thrift, Avro, and Protobuf. It will be exploited to optimize queries over massive data sets, much faster than the older file formats. Similarly, Impala is purpose built optimized query engine (that relies on Parquet).

Machine Learning and  
other advanced analytics  
are hard to write  
and perform poorly  
on Hadoop...

... but alternatives are emerging, including Spark, Storm, and proprietary solutions.

# A Final Emerging Trend...



Sunday, June 9, 13

Both are examples of technologies focused on particular problems.  
photo: Trump hotel and residences on the Chicago River.

# *Probabilistic Programming*

- Languages for Probabilistic Graphical Models??
- Bayesian Networks.
- Markov Chains.
- 
- • •

# Final Thoughts



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

Photo: Gull on a pier near Fort Mason, SF

# Consider Small-batch, Artisanal Data...

If you have lots of data, it can be valuable, but you also have lots of work to do to manage and exploit it. Finding the right amount of “curated” data to use and a matching architecture still have valuable benefits, as always. Unfortunately, I see people jump on the Big Data bandwagon who think there’s gold in their few GBs of data... That said, the “medium data” market is not well served at the moment. Hadoop can be overkill, yet traditional tools can be too expensive or not handle the load...

# SQL Strikes Back!



Copyright © 2011-2013, Dean Wampler, All Rights Reserved

Sunday, June 9, 13

NoSQL solves meets lots of requirements better than traditional RDBMSs, but people loves them some SQL!!

# Don't overlook SQL

- It's entrenched.
- Organizations want a SQL solution if they can have it.

# Hadoop owes a lot of its popularity to Hive!

Some “NoSQL”  
databases have added  
query languages (e.g.,  
Cassandra, MongoDB).

“NewSQL” databases  
are bringing NoSQL  
performance to the  
relational model.

# Examples

- Google Spanner and F1.
- NuoDB.
- VoltDB.

Spanner is the successor to BigTable. It is a globally-distributed database (consistency is maintained using the Paxos algo. and hardware synchronized clocks through GPS and atomic clocks!) Each table requires a primary key. F1 is an RDBMS built on top of it.

NuoDB is a cloud based RDBMS.

VoltDB is an example “in-memory” database, which are ideal for lots of small transactions that leverage indexing and rarely require full table scans.

# What does this mean for NoSQL?

- *People love SQL...*
- *but NoSQL meets other requirements.*

# Questions?



NoSQL Search Roadshow  
San Francisco, June 6, 2013  
[dean@concurrentthought.com](mailto:dean@concurrentthought.com)  
[@deanwampler](https://twitter.com/deanwampler)  
[polyglotprogramming.com/talks](http://polyglotprogramming.com/talks)



Sunday, June 9, 13

Copyright © 2011-2013, Dean Wampler, All Rights Reserved

All pictures Copyright © Dean Wampler, 2011-2013, All Rights Reserved. All other content is free to use, but attribution is requested.  
Photo: Same sunrise, in Burlingame.