



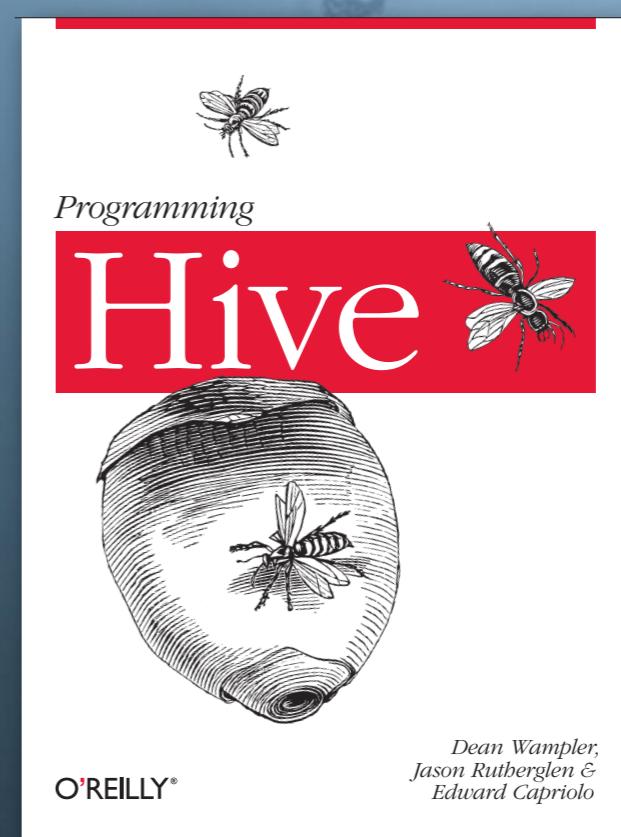
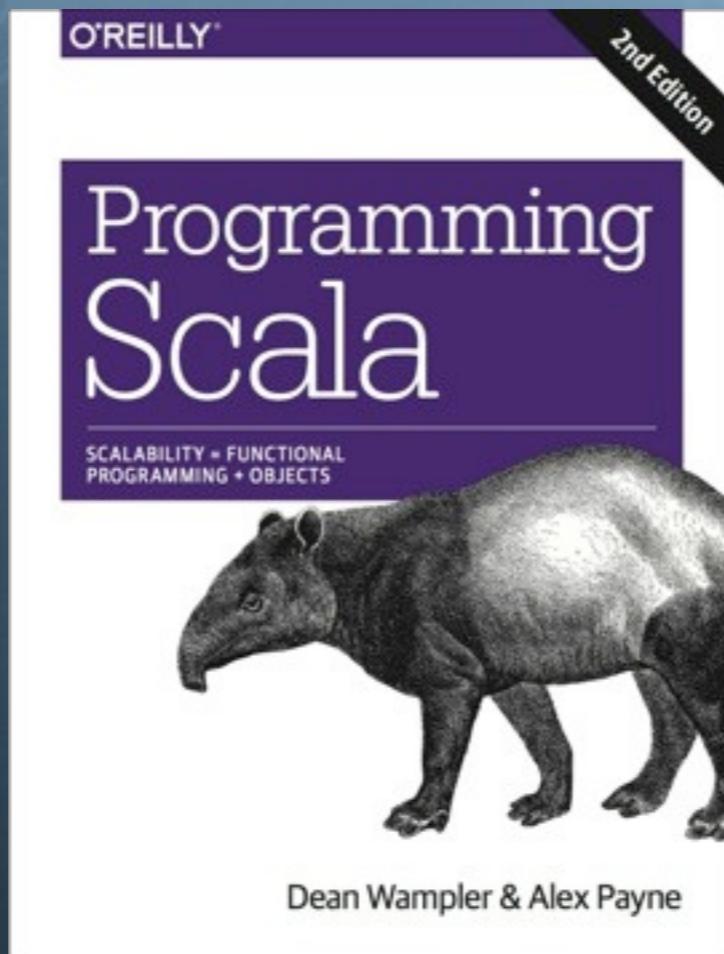
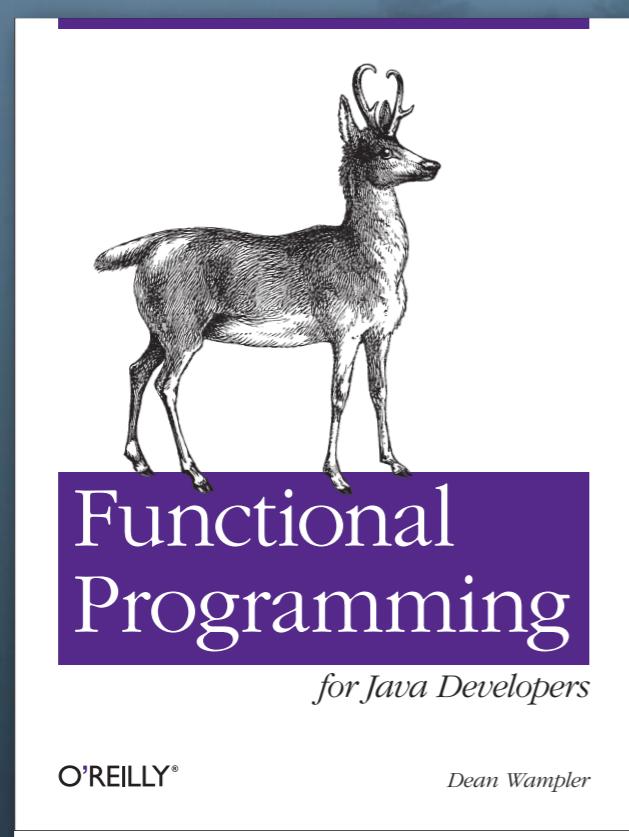
# BIG DATA STATE OF THE ART: SPARK AND THE SQL RESURGENCE

Dean Wampler, Ph.D.  
*Typesafe*

INTERNATIONAL  
SOFTWARE DEVELOPMENT  
CONFERENCE

Tuesday, September 30, 14

# Dean Wampler



[dean.wampler@typesafe.com](mailto:dean.wampler@typesafe.com)  
[polyglotprogramming.com/talks](http://polyglotprogramming.com/talks)  
[@deanwampler](https://twitter.com/deanwampler)

2

Tuesday, September 30, 14

About me. You can find this presentation and others on Big Data and Scala at [polyglotprogramming.com](http://polyglotprogramming.com).

# It's 2014



3

Tuesday, September 30, 14

A scenic mountain landscape featuring a paved path leading to a viewpoint. A person in a blue jacket and hat walks away from the camera on the path. The background shows rolling hills and mountains under a sky filled with white and grey clouds.

Hadoop has been  
very successful.



But it's not perfect

5



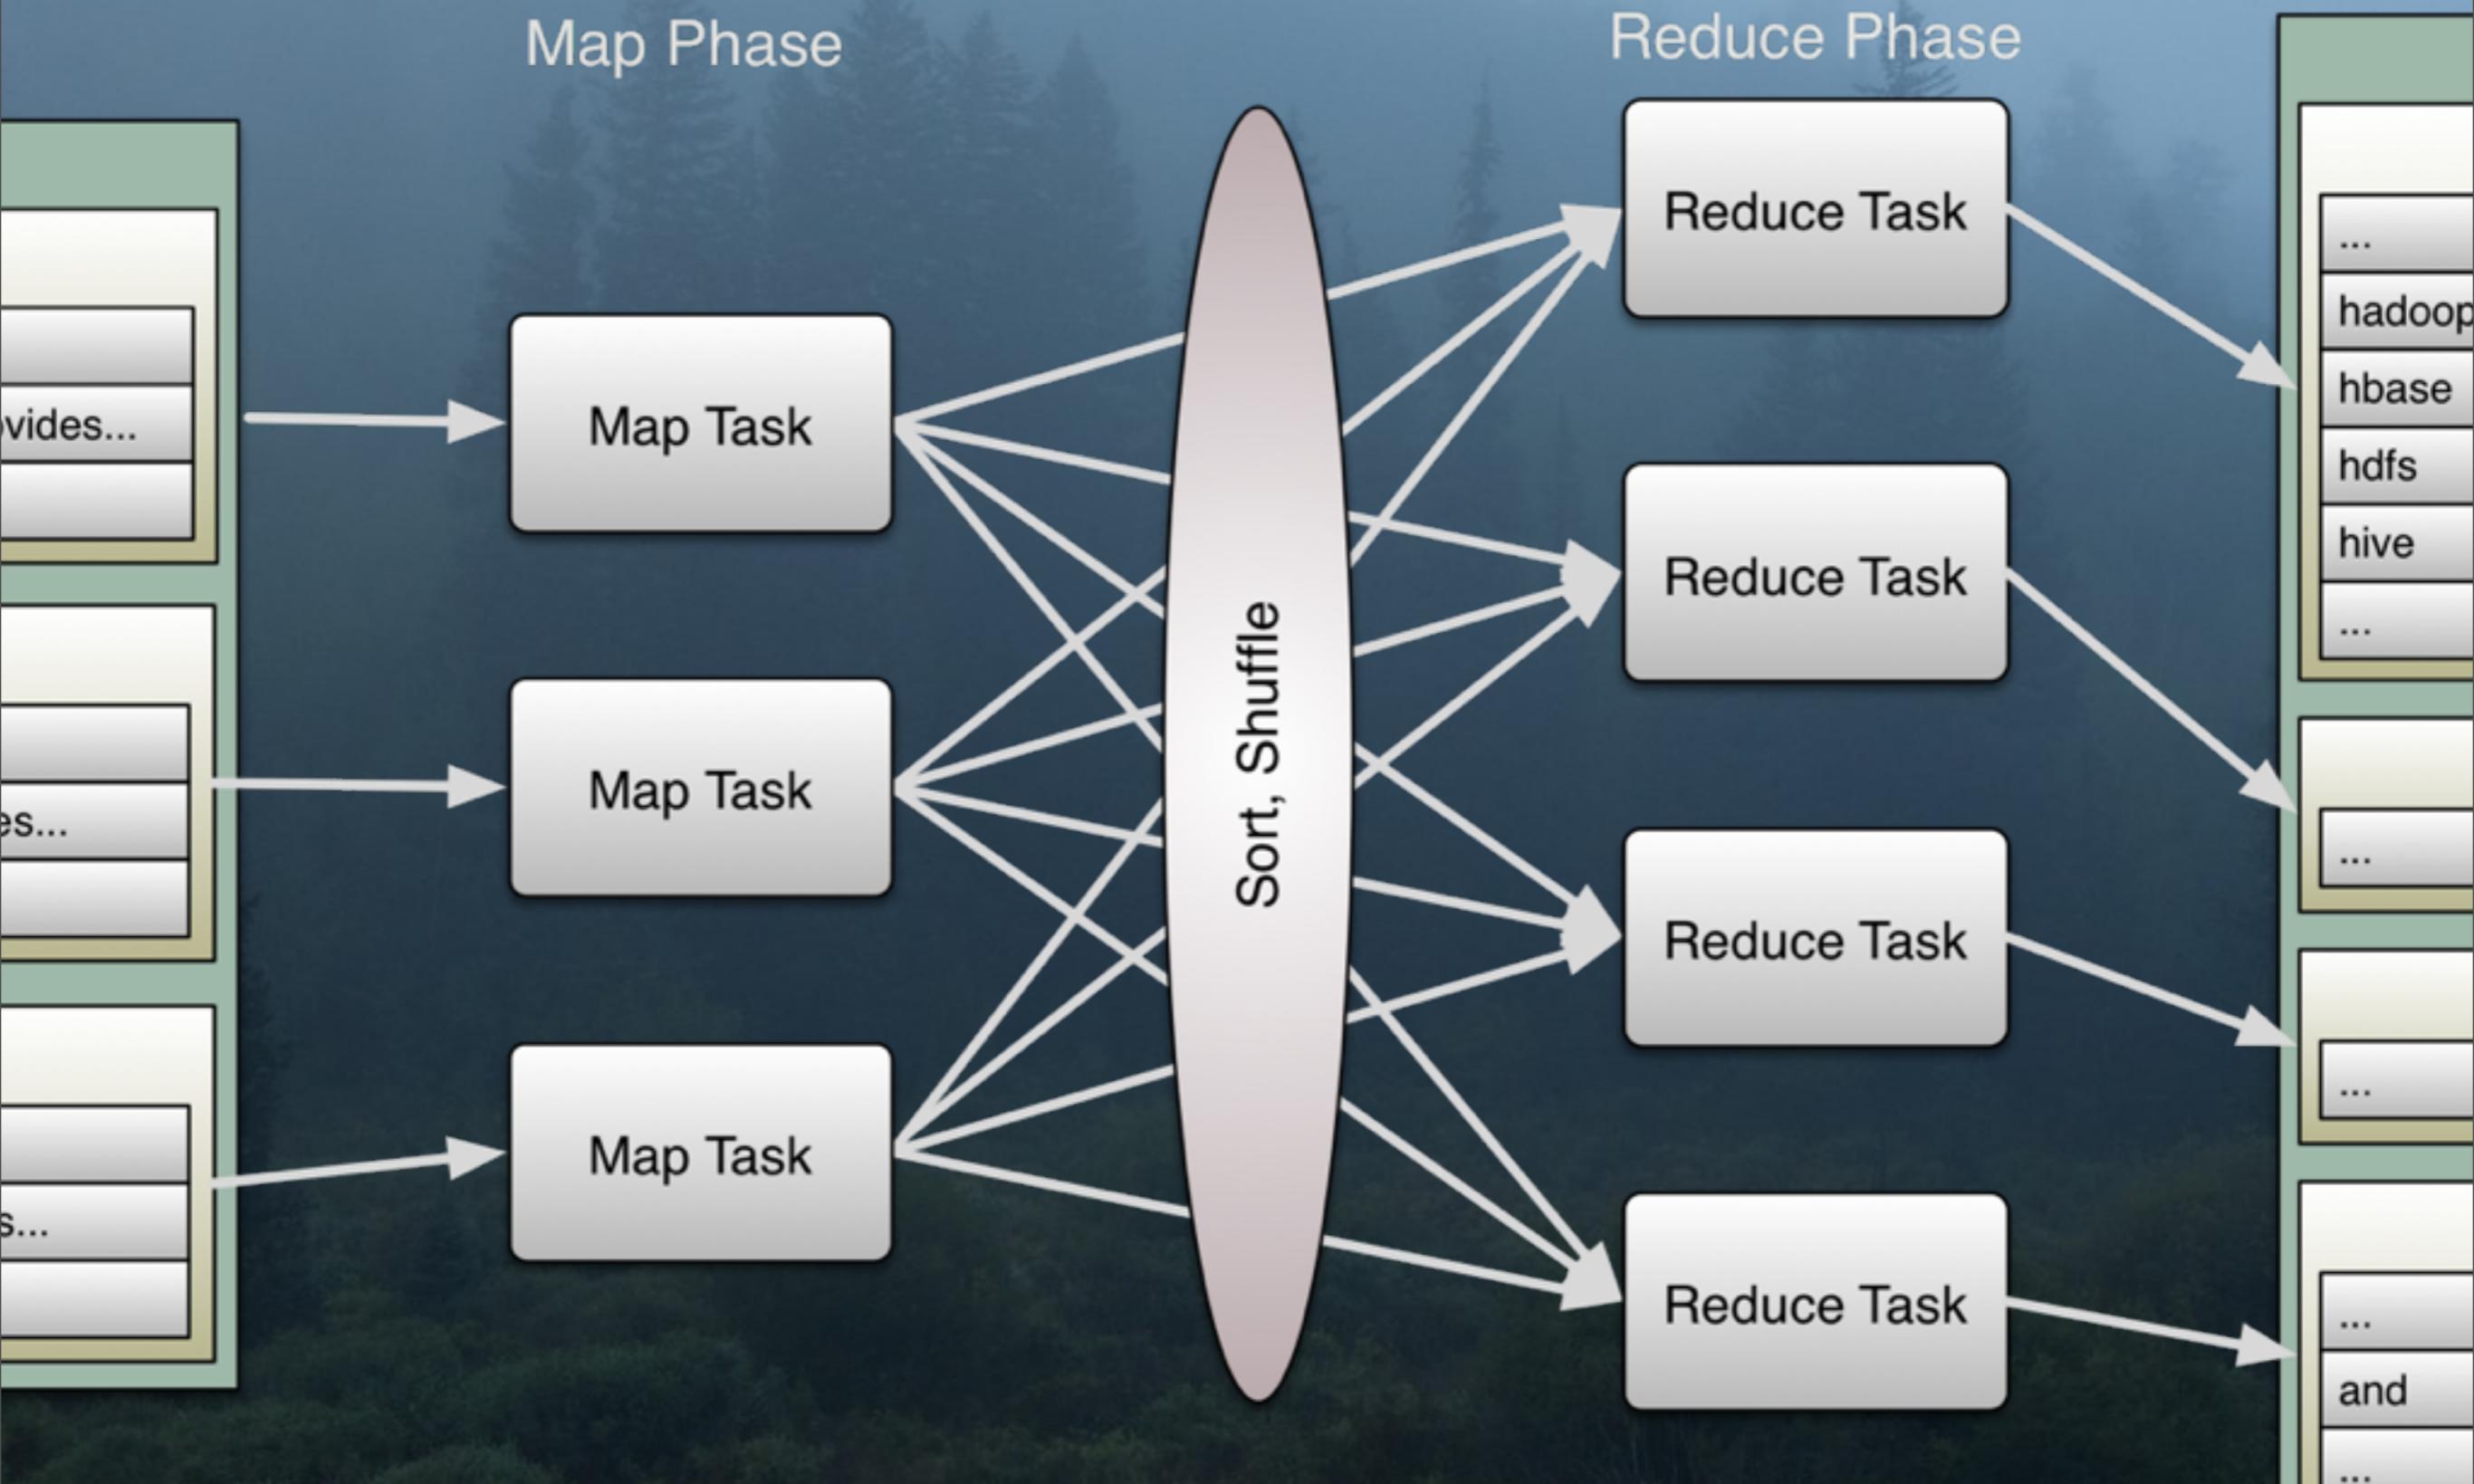
# MapReduce

6

Tuesday, September 30, 14

The limitations of MapReduce have become increasingly significant...

# 1 Map step + 1 Reduce step



Tuesday, September 30, 14

You get one map step and one reduce step. You can sequence jobs together when necessary.

# Problems



Limited  
programming  
model

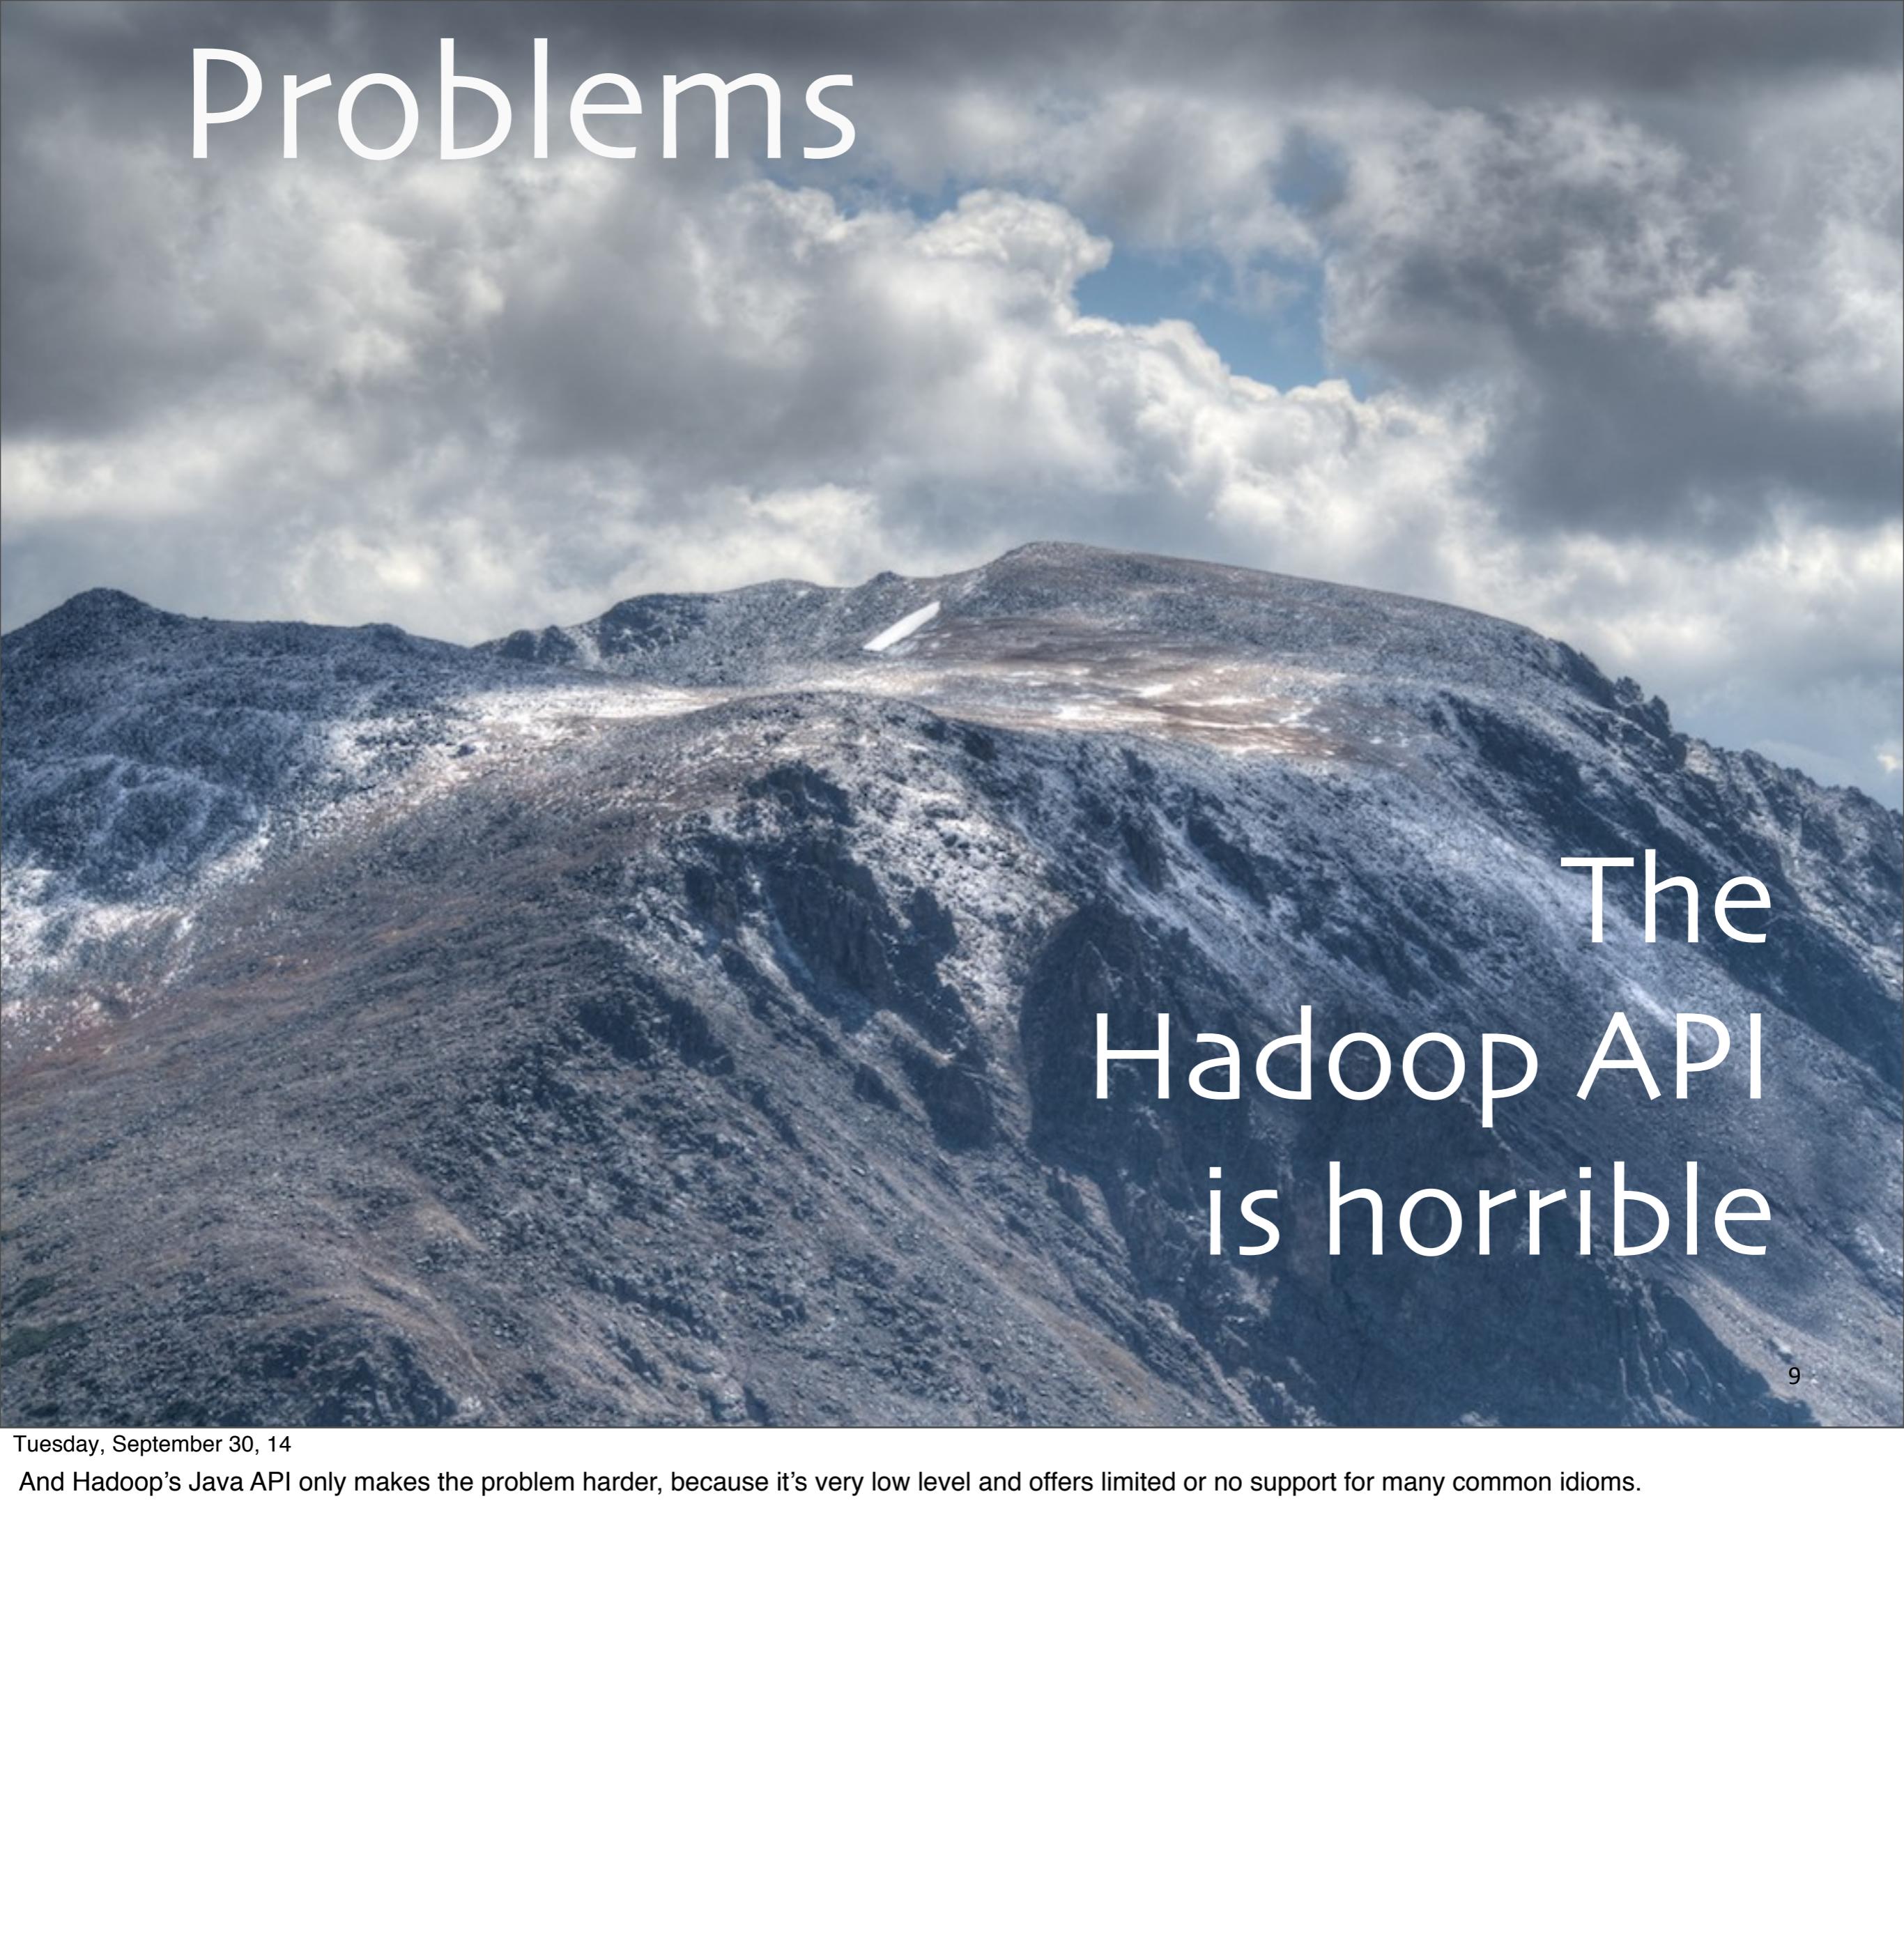
8

Tuesday, September 30, 14

MapReduce is a restrictive model. Writing jobs requires arcane, specialized skills that few master. It's not easy mapping arbitrary algorithms to this model. For example, algorithms that are naturally iterative are especially hard, because MR doesn't support iteration efficiently. For a good overview, see <http://lintool.github.io/MapReduceAlgorithms/>.

The limited model doesn't just impede developer productivity, it makes it much harder to build tools on top of the model, as we'll discuss.

# Problems



The  
Hadoop API  
is horrible

9

Tuesday, September 30, 14

And Hadoop's Java API only makes the problem harder, because it's very low level and offers limited or no support for many common idioms.

# Example

Inverted  
Index

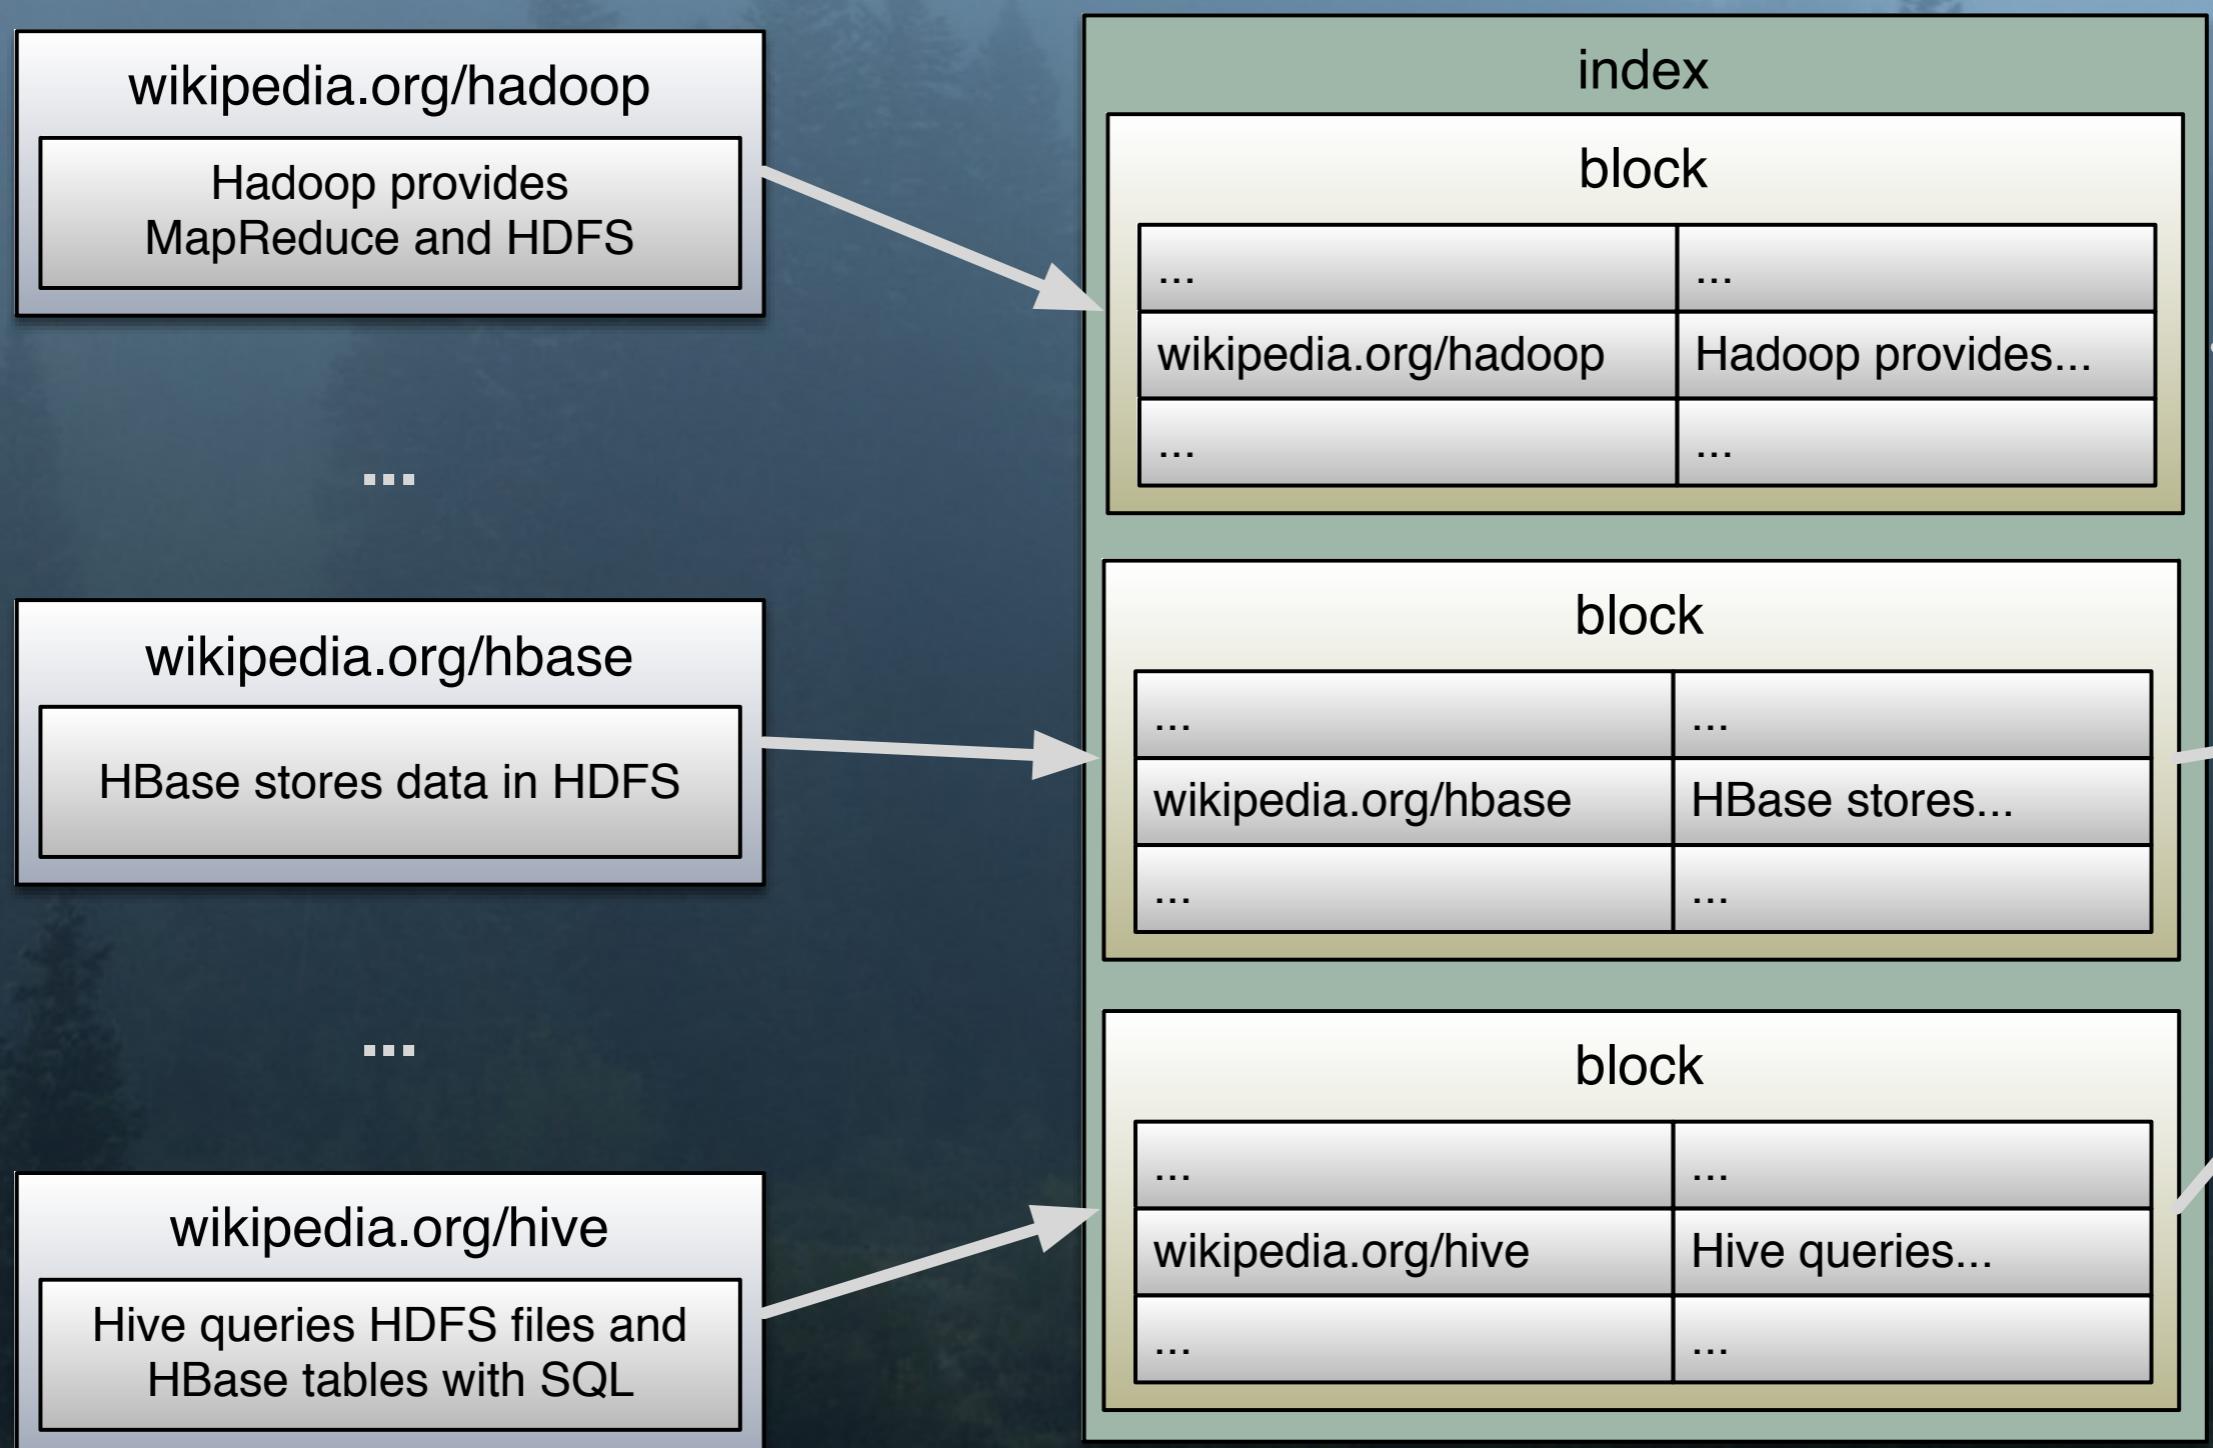
10

Tuesday, September 30, 14

See compare and contrast MR with Spark, let's use this classic algorithm.

# Inverted Index

Web Crawl



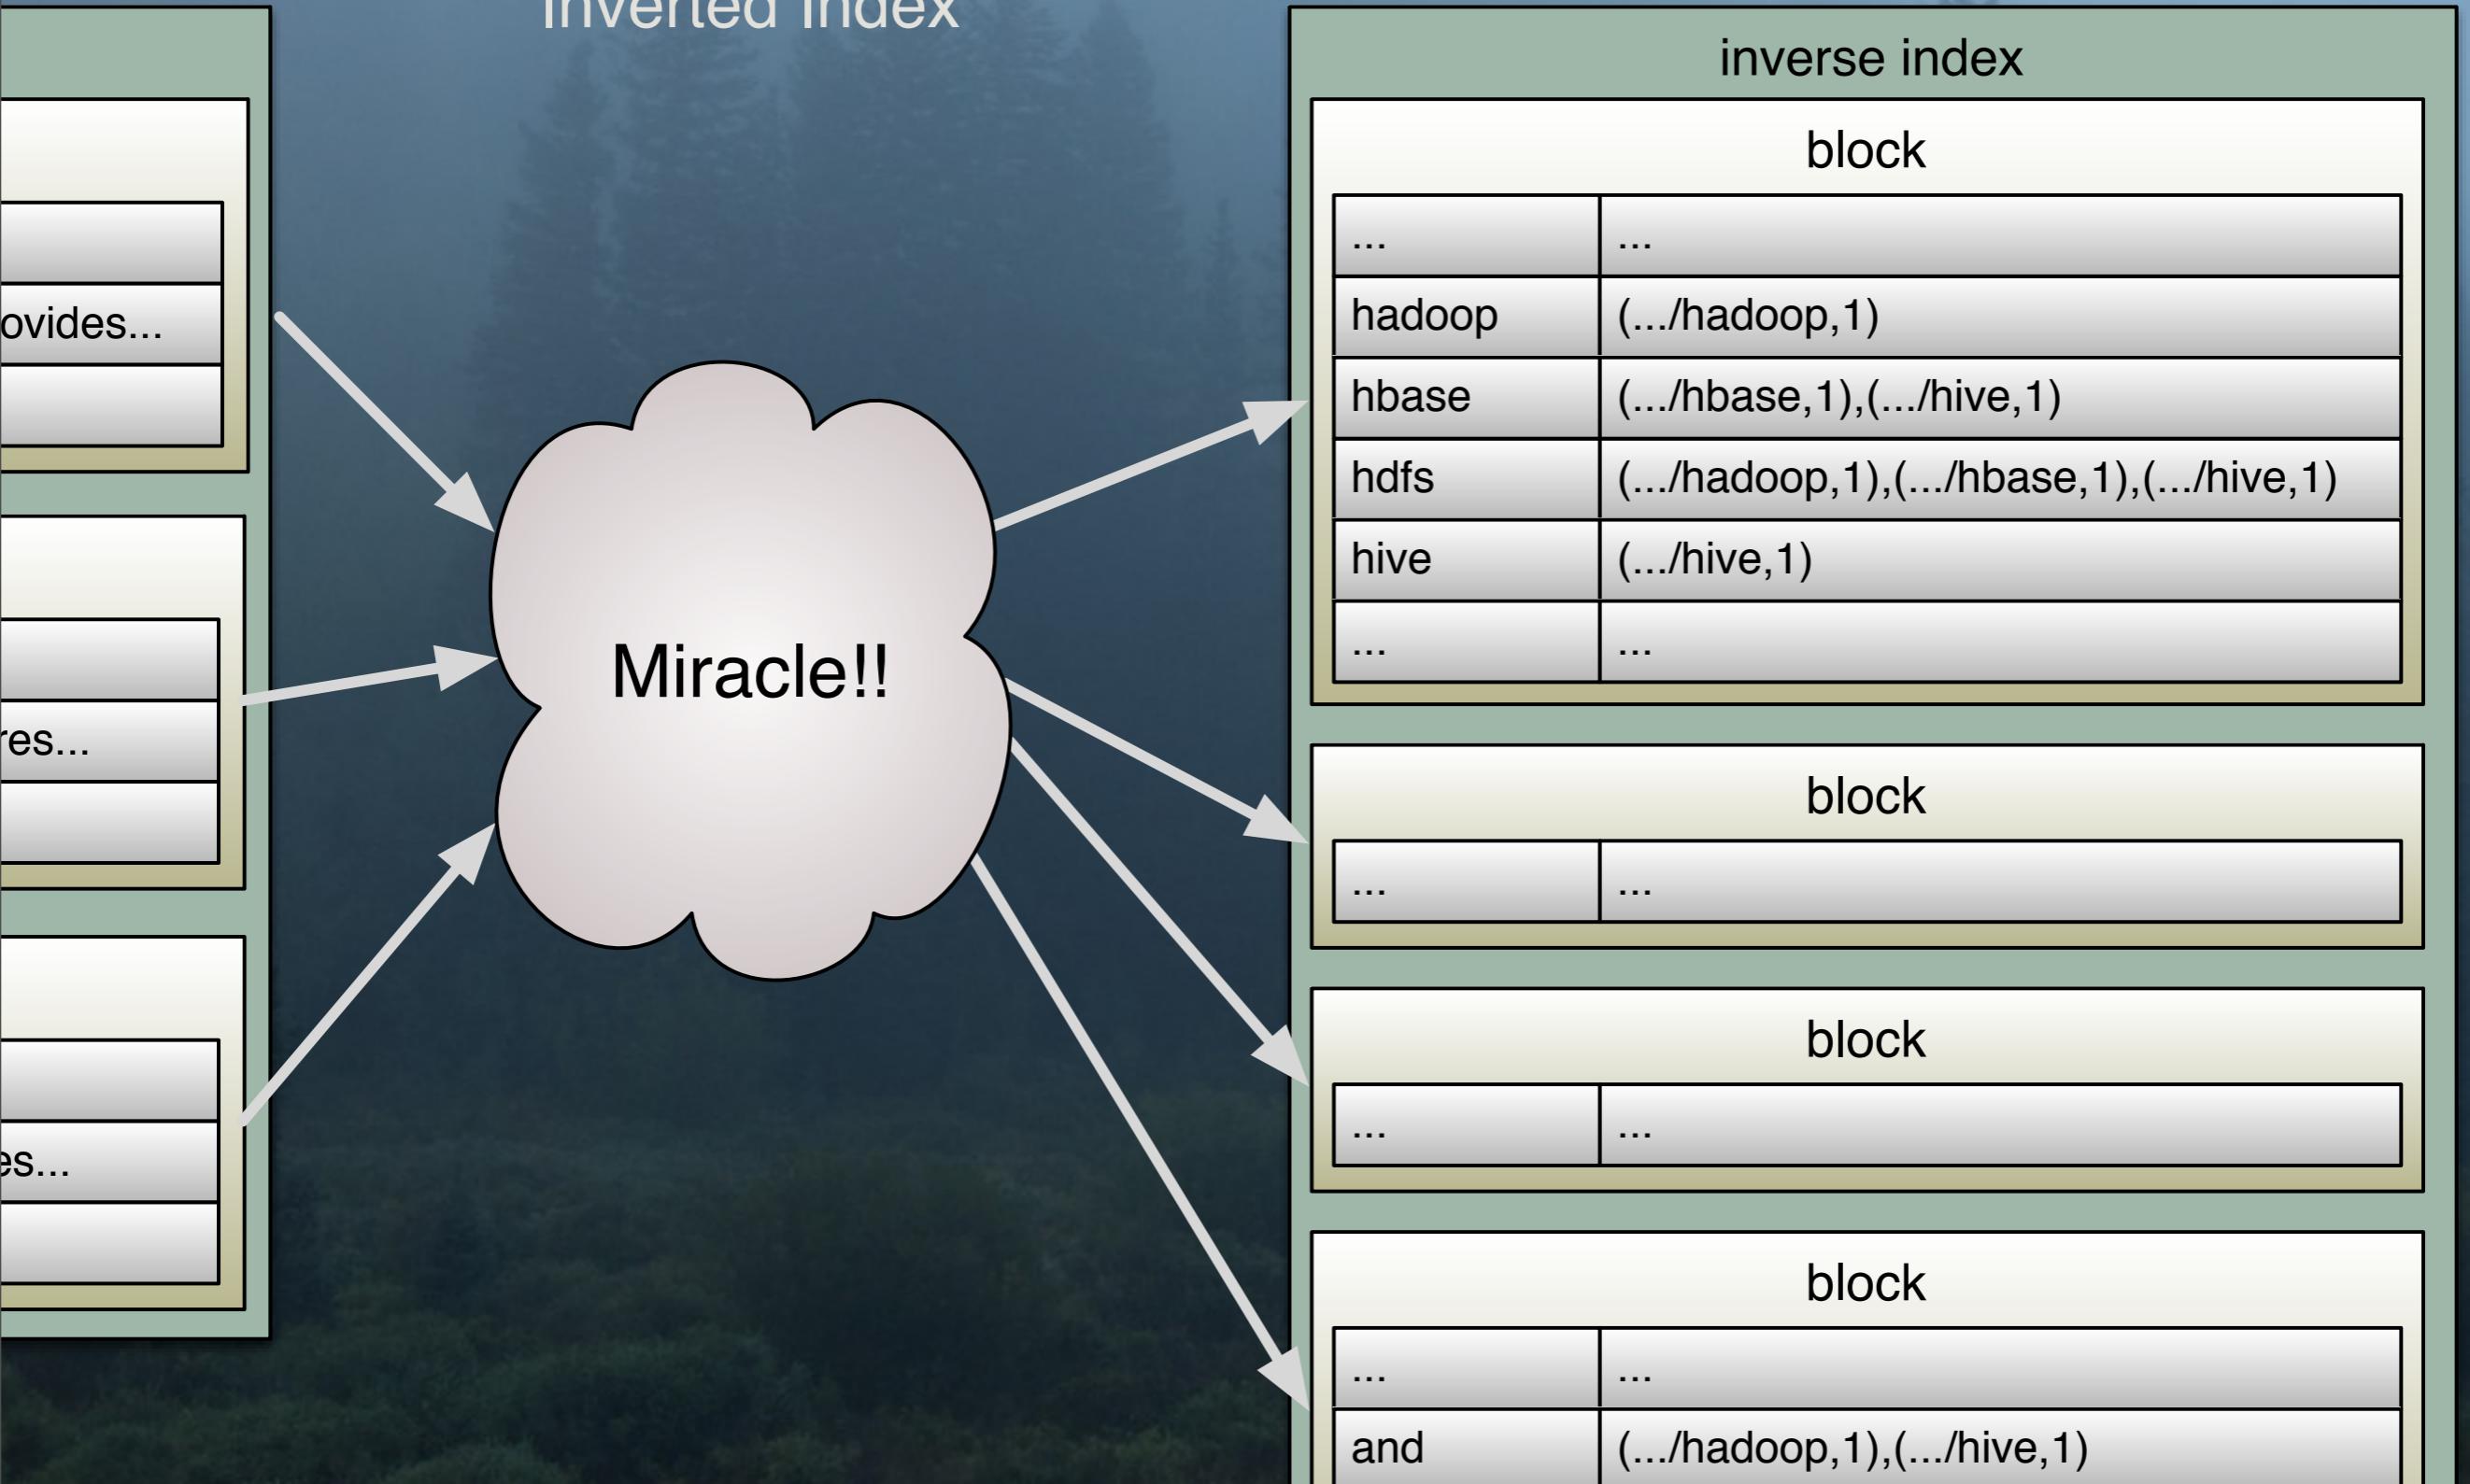
11

Tuesday, September 30, 14

First we crawl the web to build a data set of document names/ids and their contents. Then we “invert” it; we tokenize the contents into words and build a new index from each word to the list of documents that contain the word and the count in each document. This is a basic data set for search engines.

# Inverted Index

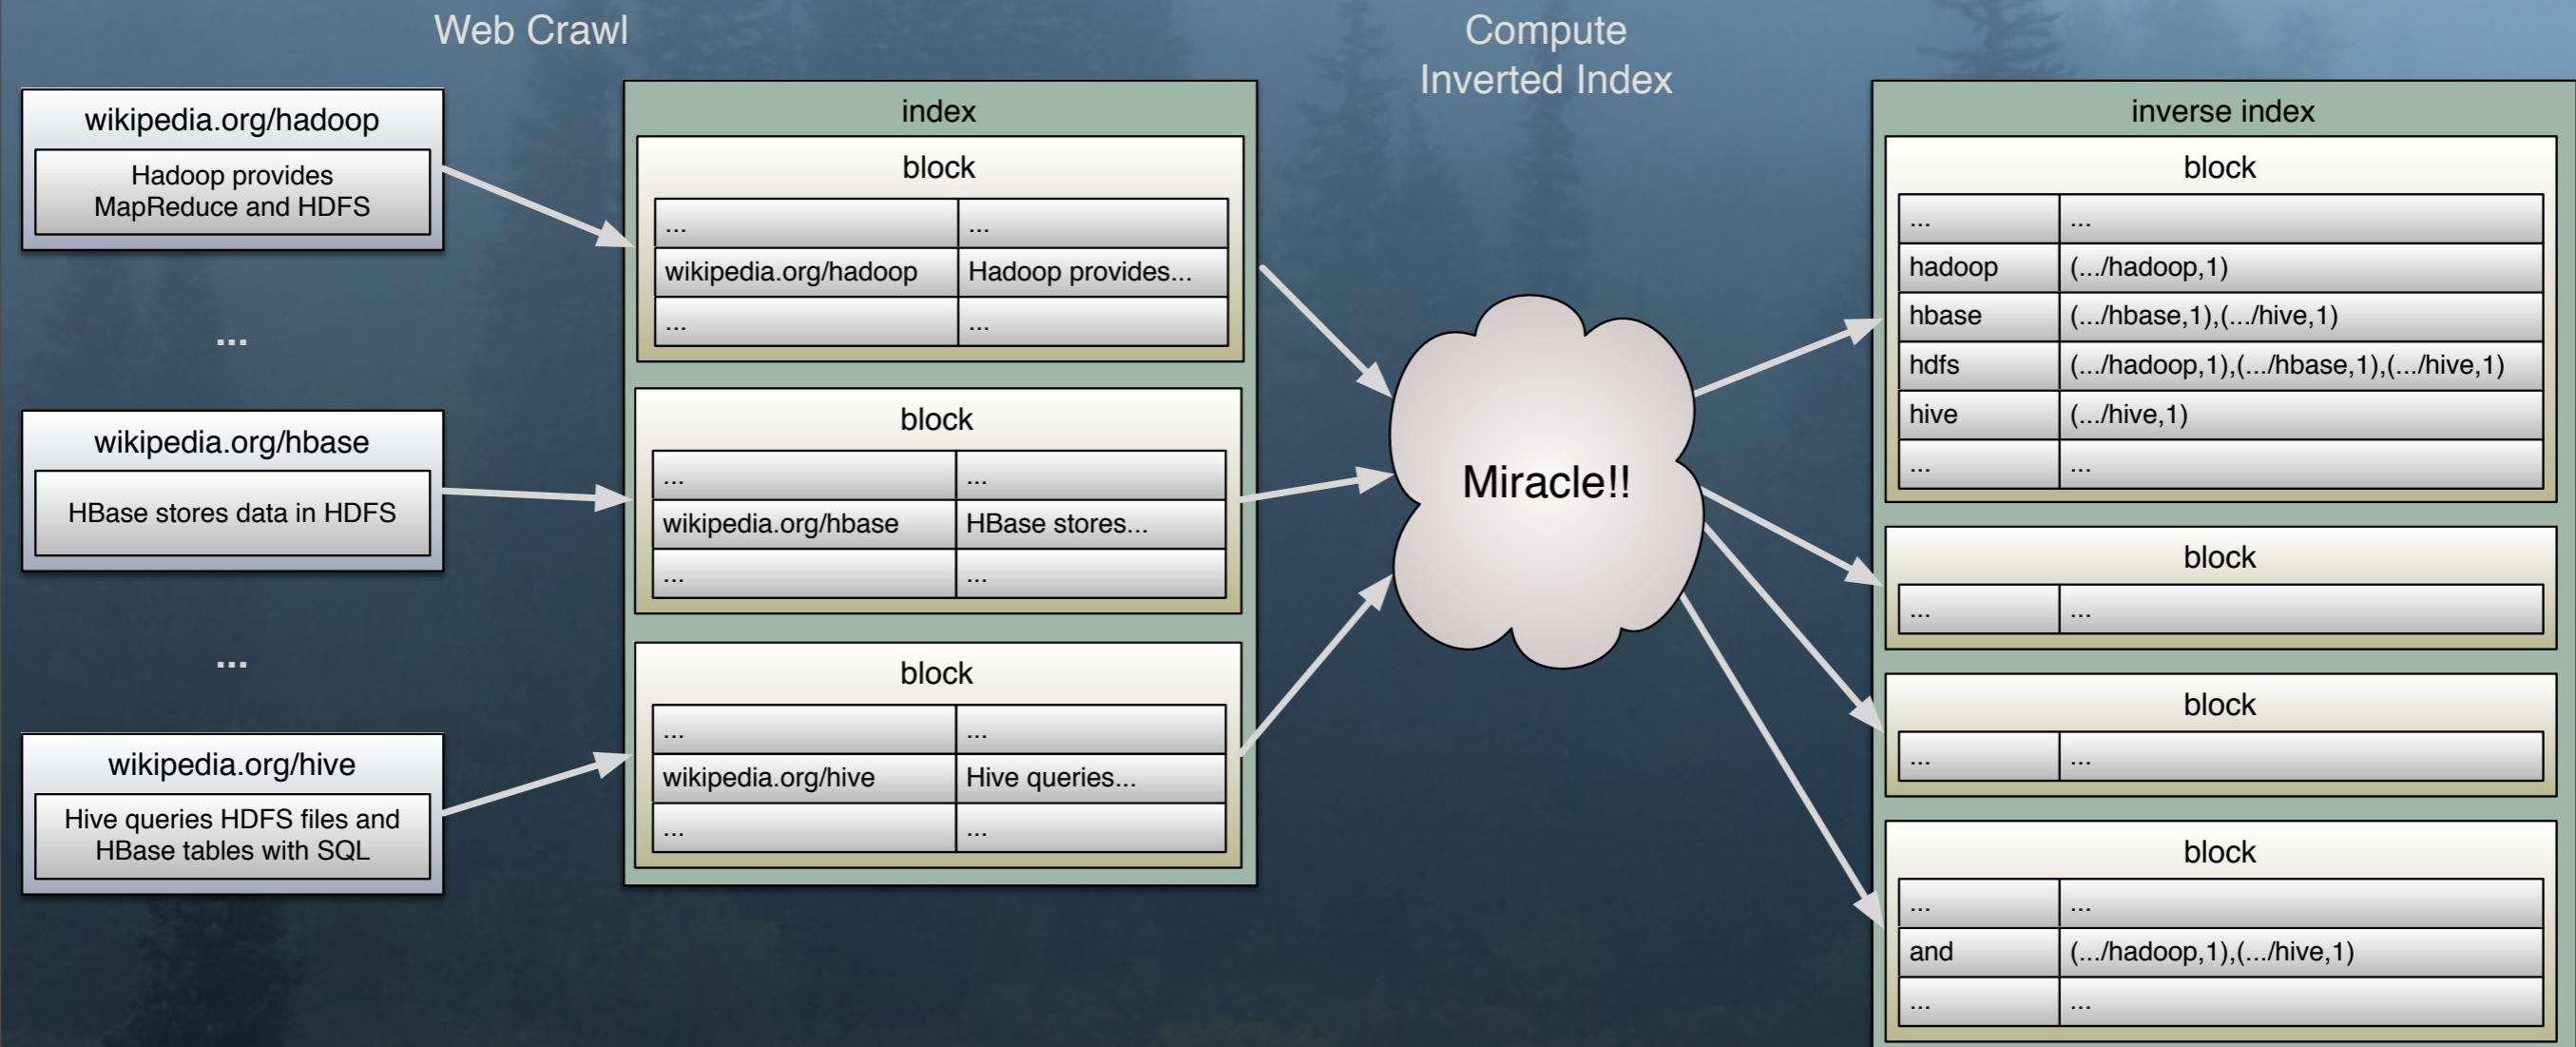
Compute  
Inverted Index



Tuesday, September 30, 14

First we crawl the web to build a data set of document names/ids and their contents. Then we “invert” it; we tokenize the contents into words and build a new index from each word to the list of documents that contain the word and the count in each document. This is a basic data set for search engines.

# Inverted Index



# Altogether

13

Tuesday, September 30, 14

We'll implement the "miracle".

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf =
            new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
```

14

Tuesday, September 30, 14

I'm not going to explain many of the details. The point is to notice all the boilerplate that obscures the problem logic.

Everything is in one outer class. We start with a main routine that sets up the job.

I used yellow for method calls, because methods do the real work!! But notice that most of the functions in this code don't really do a whole lot of work for us...

```
JobClient client = new JobClient();
JobConf conf =
    new JobConf(LineIndexer.class);

conf.setJobName("LineIndexer");
conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(Text.class);
FileInputFormat.addInputPath(conf,
    new Path("input"));
FileOutputFormat.setOutputPath(conf,
    new Path("output"));
conf.setMapperClass(
    LineIndexMapper.class);
conf.setReducerClass(
    LineIndexReducer.class);

client.setConf(conf);
```

15

```
    LineIndexMapper.class);  
    conf.setReducerClass(  
        LineIndexReducer.class);  
  
    client.setConf(conf);  
  
    try {  
        JobClient.runJob(conf);  
    } catch (Exception e) {  
        e.printStackTrace();  
    }  
}  
  
public static class LineIndexMapper  
    extends MapReduceBase  
    implements Mapper<LongWritable, Text,  
              Text, Text> {
```

16

Tuesday, September 30, 14

main ends with a try-catch clause to run the job.

```
public static class LineIndexMapper  
extends MapReduceBase  
implements Mapper<LongWritable, Text,  
Text, Text> {  
private final static Text word =  
new Text();  
private final static Text location =  
new Text();  
  
public void map(  
LongWritable key, Text val,  
OutputCollector<Text, Text> output,  
Reporter reporter) throws IOException {  
  
FileSplit fileSplit =  
(FileSplit)reporter.getInputSplit();  
String fileName -
```

17

Tuesday, September 30, 14

This is the LineIndexMapper class for the mapper. The map method does the real work of tokenization and writing the (word, document-name) tuples.

```
FileSplit fileSplit =
  (FileSplit)reporter.getInputSplit();
String fileName =
  fileSplit.getPath().getName();
location.set(fileName);

String line = val.toString();
StringTokenizer itr = new
  StringTokenizer(line.toLowerCase());
while (itr.hasMoreTokens()) {
  word.set(itr.nextToken());
  output.collect(word, location);
}
}
```

18

Tuesday, September 30, 14

The rest of the LineIndexMapper class and map  
method.

```
public static class LineIndexReducer
extends MapReduceBase
implements Reducer<Text, Text,
Text, Text> {
public void reduce(Text key,
Iterator<Text> values,
OutputCollector<Text, Text> output,
Reporter reporter) throws IOException {
boolean first = true;
StringBuilder toReturn =
new StringBuilder();
while (values.hasNext()) {
if (!first)
toReturn.append(", ");
first=false;
toReturn.append(
values.next().toString());
}
```

19

Tuesday, September 30, 14

The reducer class, LineIndexReducer, with the reduce method that is called for each key and a list of values for that key. The reducer is stupid; it just reformats the values collection into a long string and writes the final (word,list-string) output.

```
boolean first = true;
StringBuilder toReturn =
    new StringBuilder();
while (values.hasNext()) {
    if (!first)
        toReturn.append(", ");
    first=false;
    toReturn.append(
        values.next().toString());
}
output.collect(key,
    new Text(toReturn.toString()));
}
```

# Altogether

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf = new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(Text.class);
        FileInputFormat.addInputPath(conf,
            new Path("input"));
        FileOutputFormat.setOutputPath(conf,
            new Path("output"));
        conf.setMapperClass(
            LineIndexMapper.class);
        conf.setReducerClass(
            LineIndexReducer.class);

        client.setConf(conf);

        try {
            JobClient.runJob(conf);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    public static class LineIndexMapper
        extends MapReduceBase
        implements Mapper<LongWritable, Text,
                    Text, Text> {
        private final static Text word =
            new Text();
        private final static Text location =
            new Text();

        public void map(
            LongWritable key, Text val,
            OutputCollector<Text, Text> output,
            Reporter reporter) throws IOException {

            FileSplit fileSplit =
                (FileSplit)reporter.getInputSplit();
            String fileName =
                fileSplit.getPath().getName();
            location.set(fileName);

            String line = val.toString();
            StringTokenizer itr = new
                StringTokenizer(line.toLowerCase());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                output.collect(word, location);
            }
        }
    }

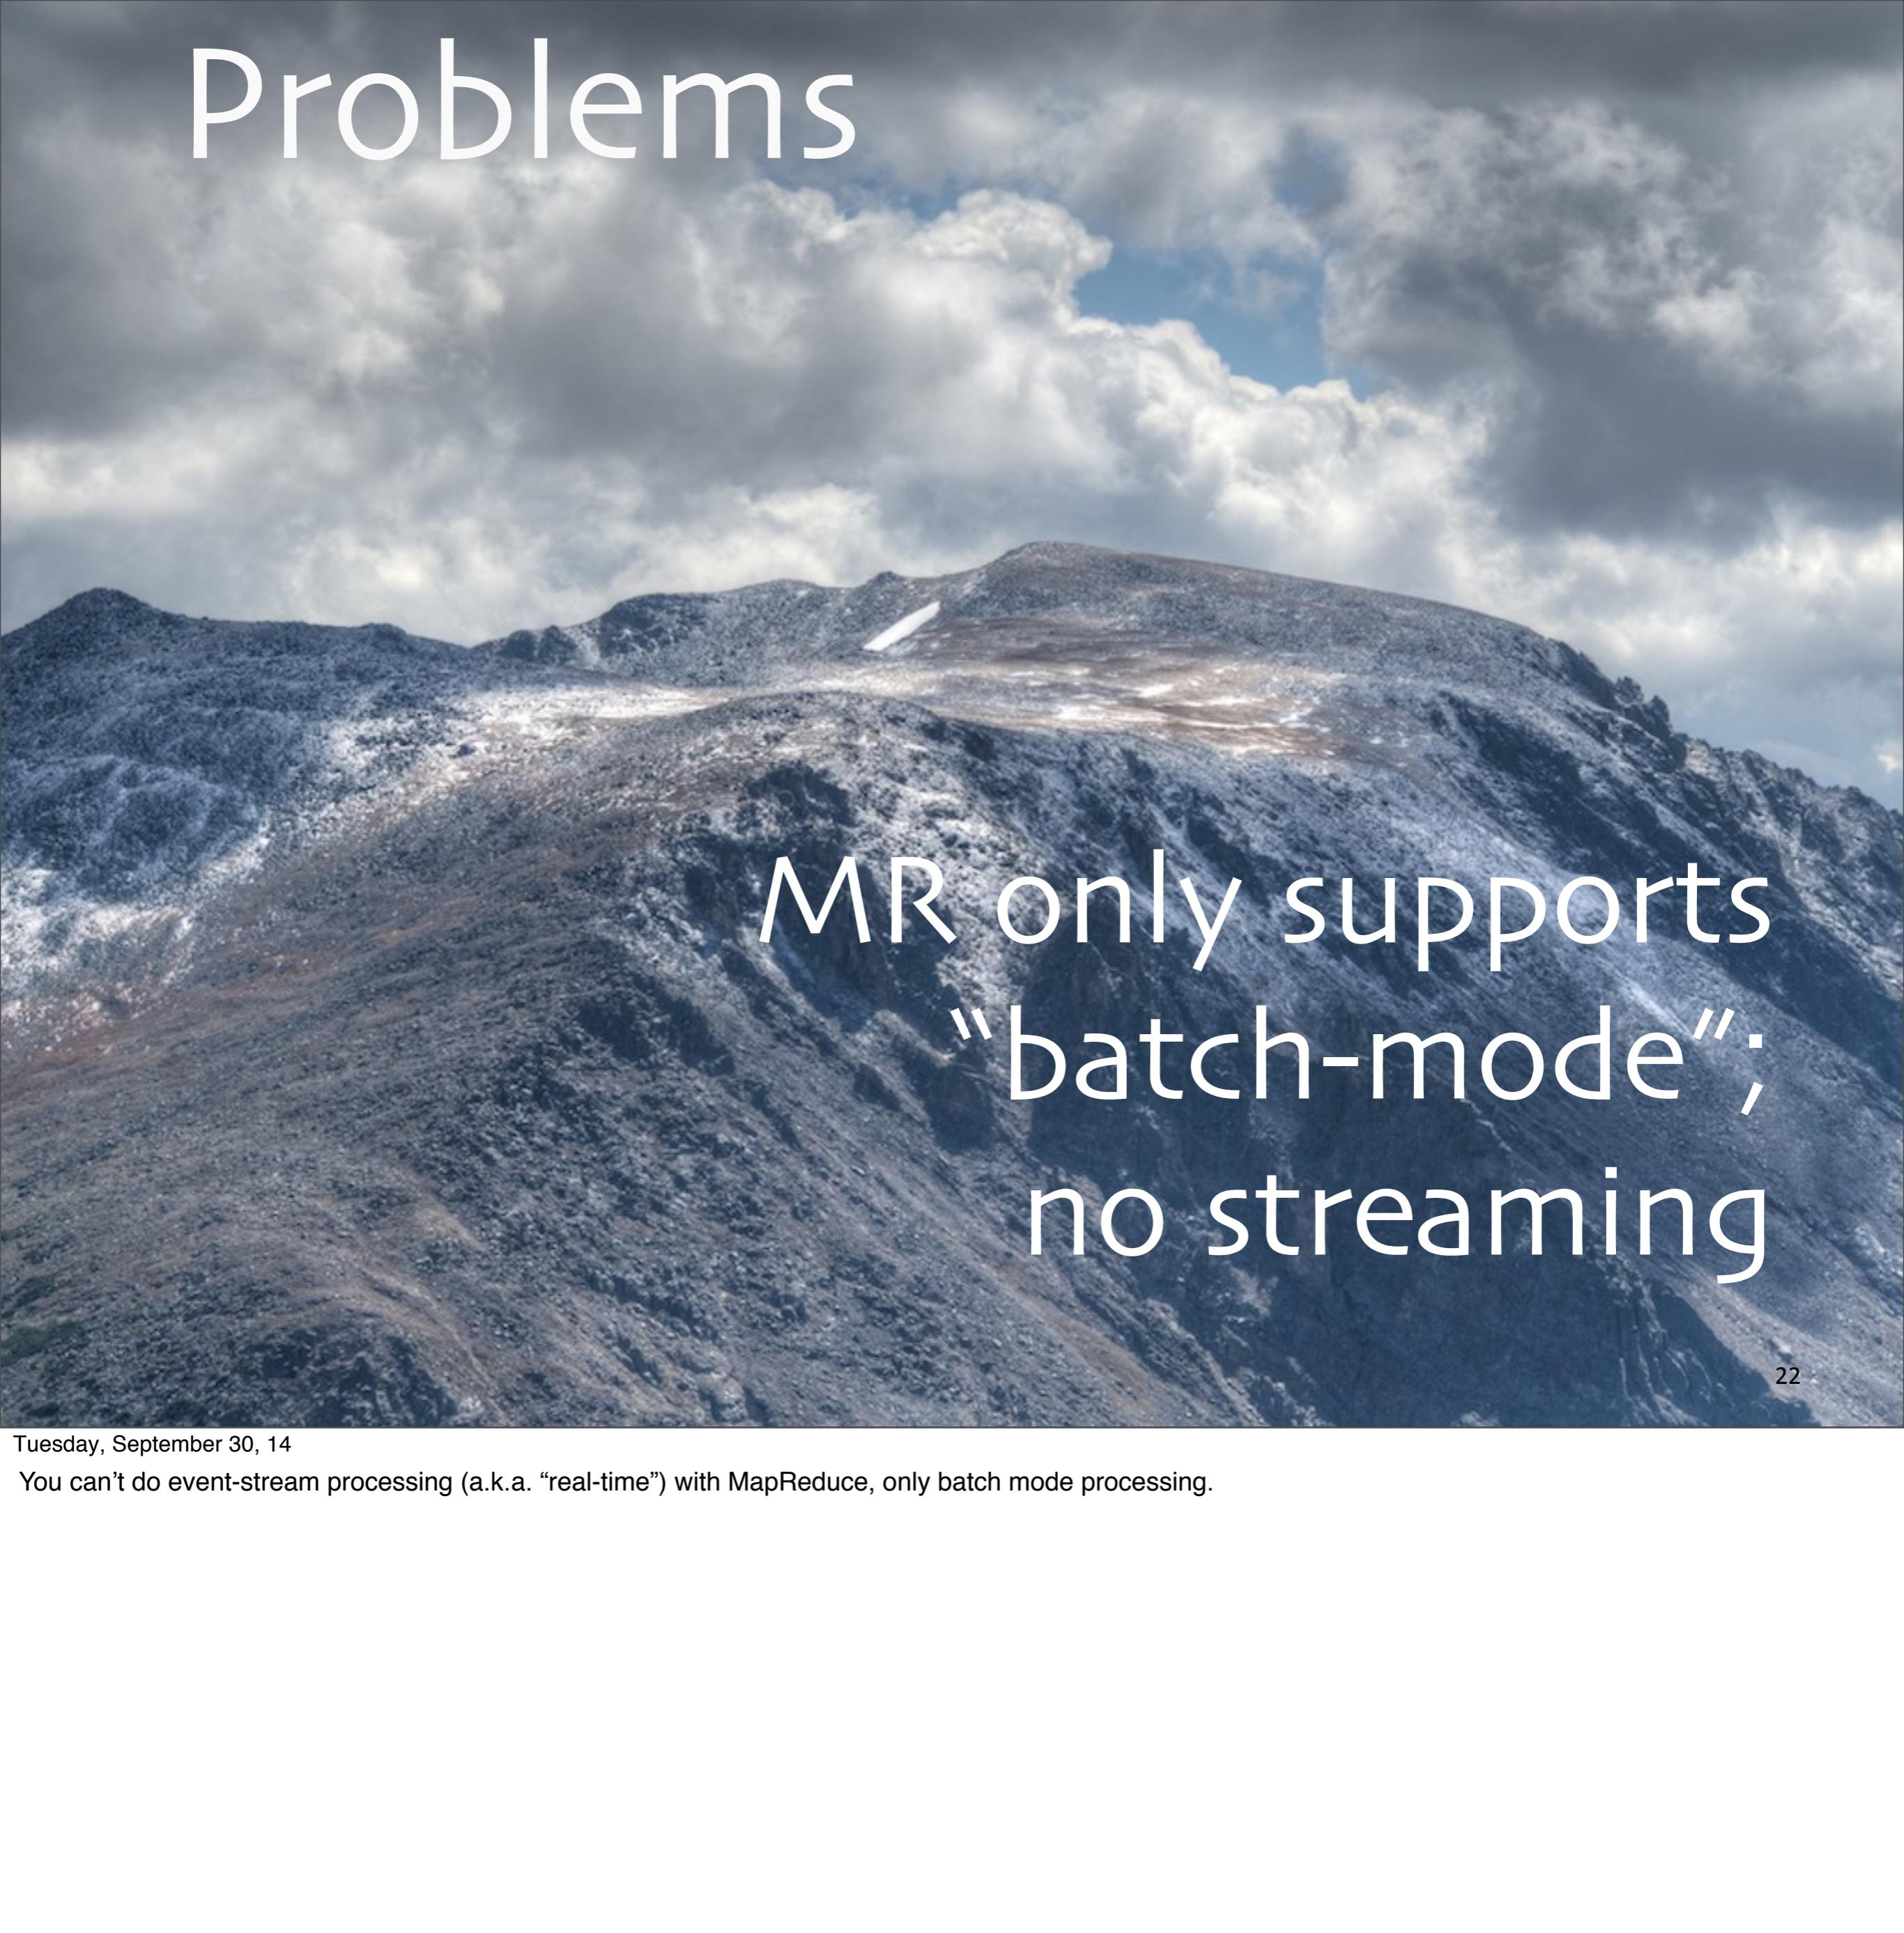
    public static class LineIndexReducer
        extends MapReduceBase
        implements Reducer<Text, Text,
                           Text, Text> {
        public void reduce(Text key,
                          Iterator<Text> values,
                          OutputCollector<Text, Text> output,
                          Reporter reporter) throws IOException {
            boolean first = true;
            StringBuilder toReturn =
                new StringBuilder();
            while (values.hasNext()) {
                if (!first)
                    toReturn.append(", ");
                first=false;
                toReturn.append(
                    values.next().toString());
            }
            output.collect(key,
                new Text(toReturn.toString()));
        }
    }
}
```

21

Tuesday, September 30, 14

The whole shebang (6pt. font) This would take a few hours to write, test, etc. assuming you already know the API and the idioms for using it.

# Problems

A wide-angle photograph of a mountain range. The mountains are dark, rocky, and covered in patches of snow or ice. The sky above is filled with heavy, grey clouds, suggesting an overcast day.

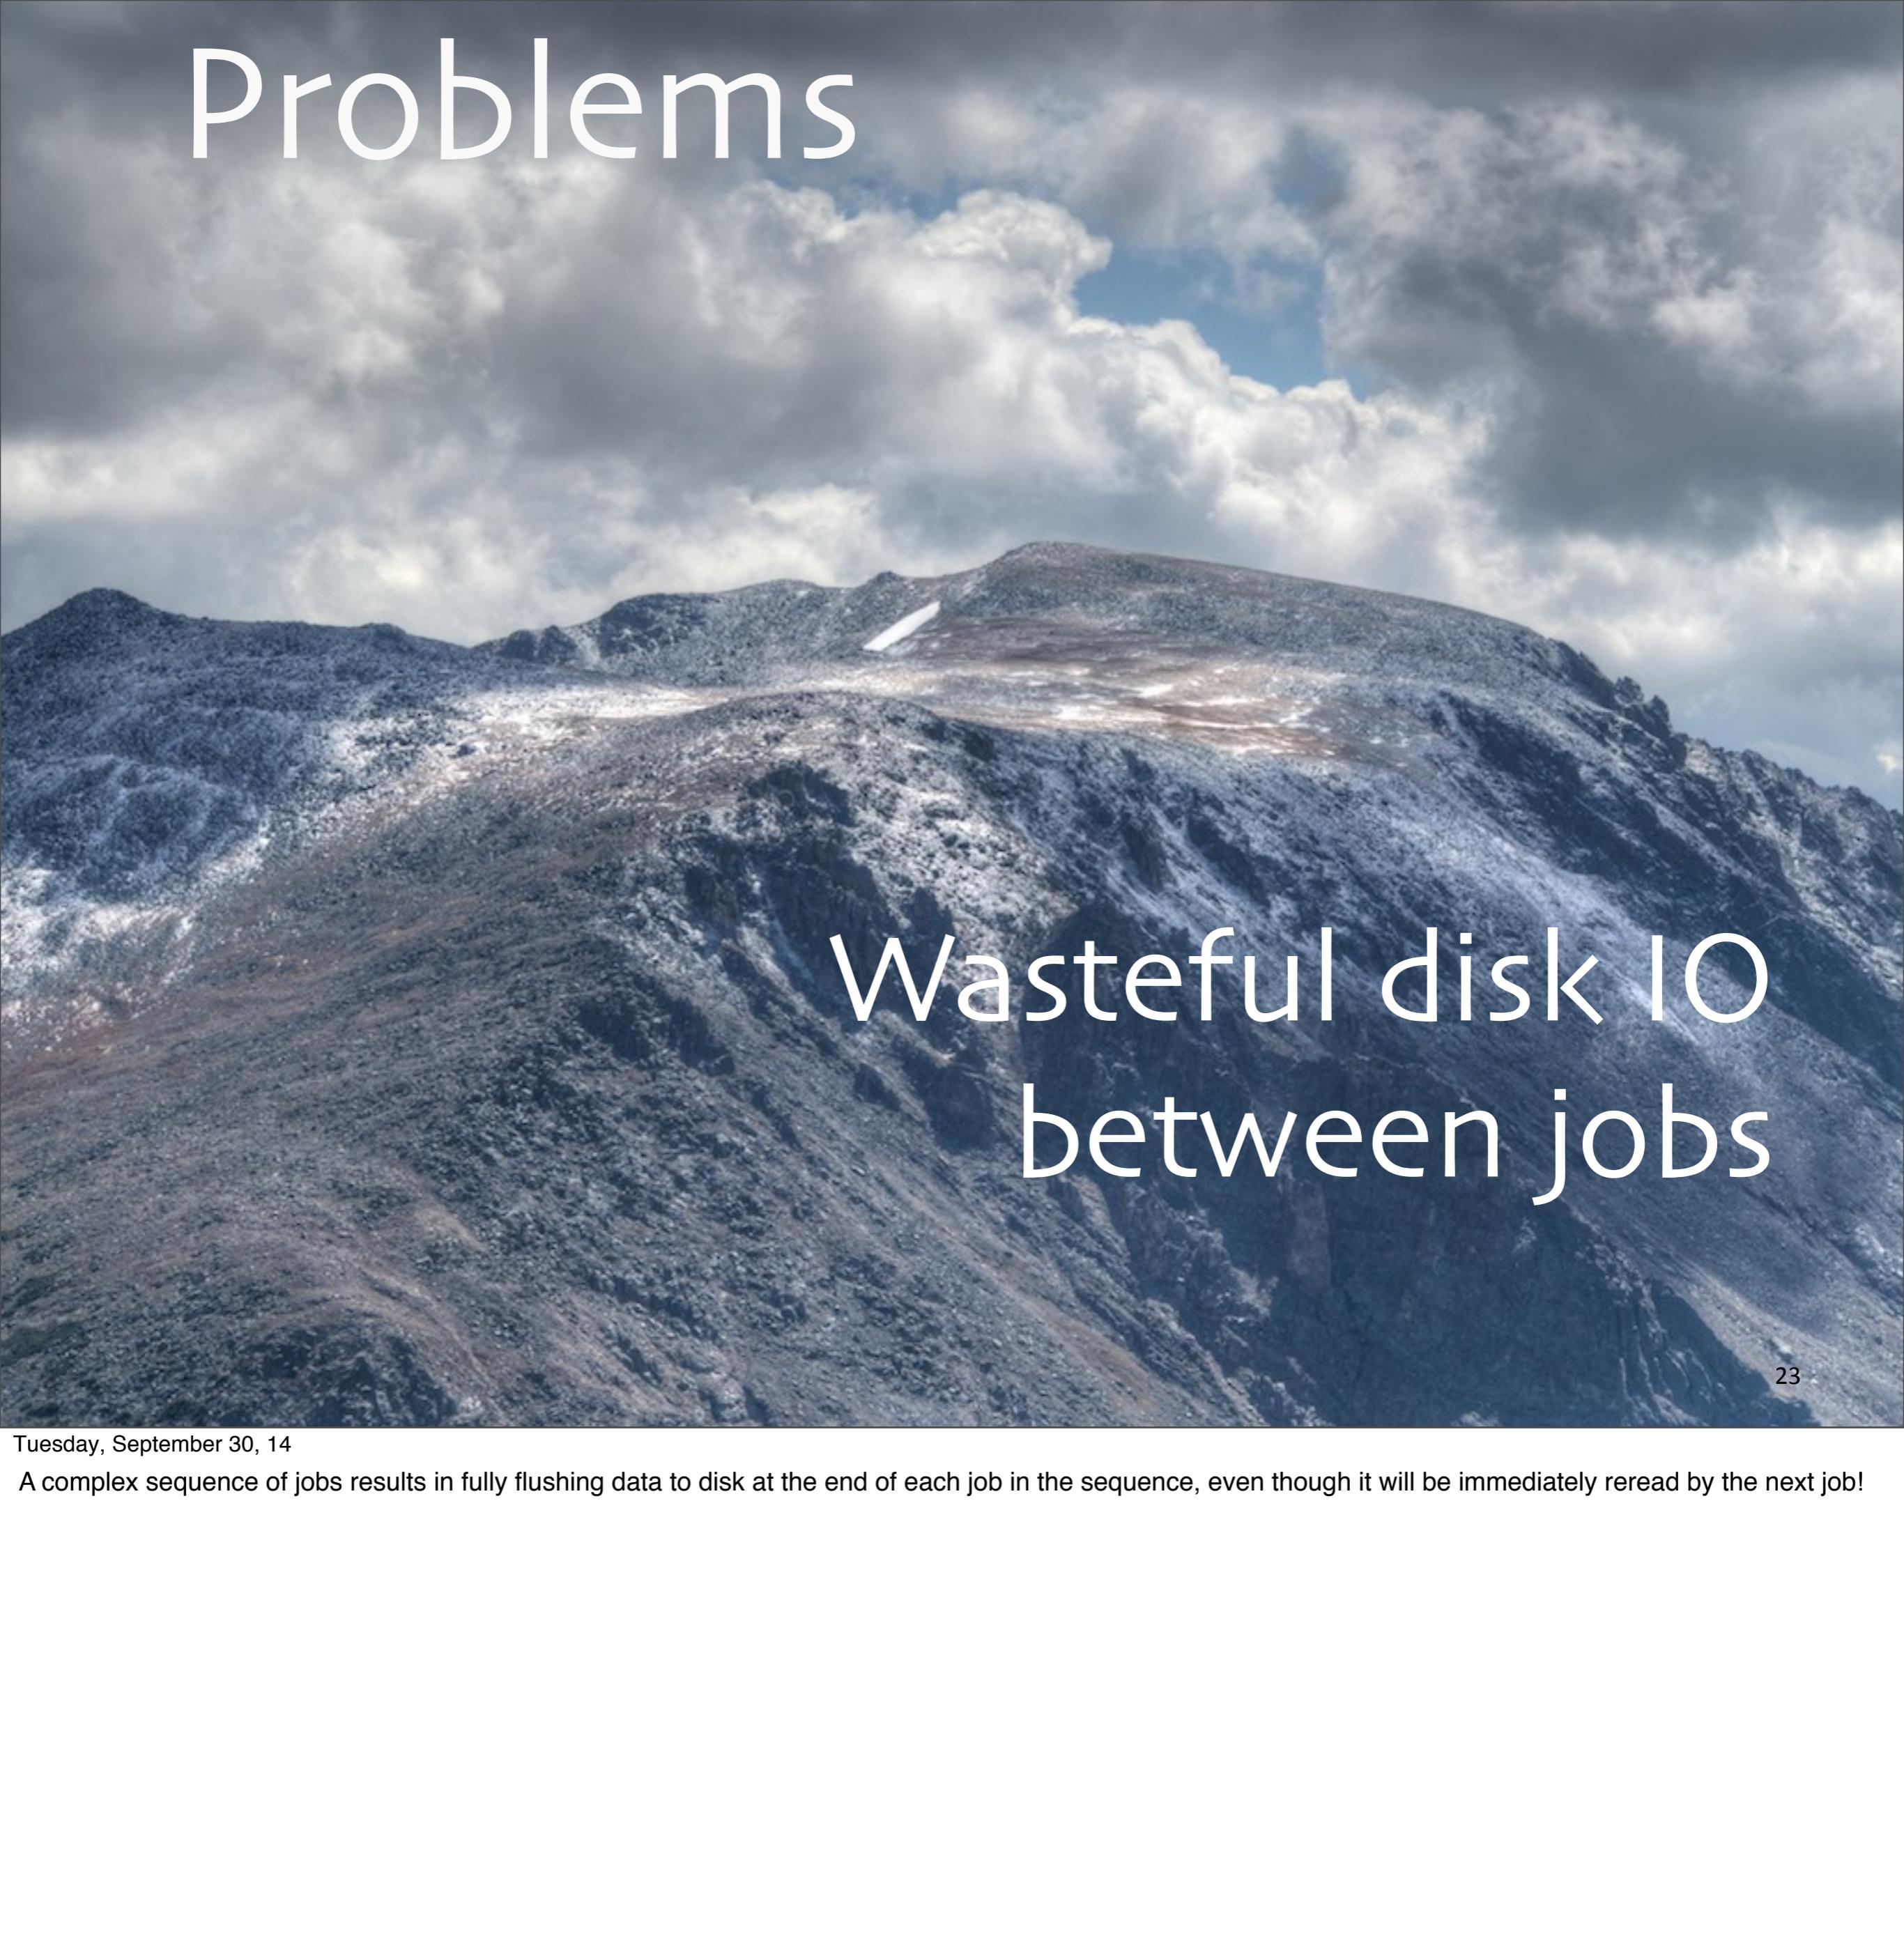
MR only supports  
“batch-mode”;  
no streaming

22

Tuesday, September 30, 14

You can't do event-stream processing (a.k.a. “real-time”) with MapReduce, only batch mode processing.

# Problems

A wide-angle photograph of a mountain range. The mountains are dark, rocky, and partially covered with snow or ice. The sky above is filled with heavy, white clouds. In the center of the image, there is a large, bright white area that appears to be a reflection or a lens flare.

Wasteful disk IO  
between jobs

23

Tuesday, September 30, 14

A complex sequence of jobs results in fully flushing data to disk at the end of each job in the sequence, even though it will be immediately reread by the next job!



# Enter Spark

24

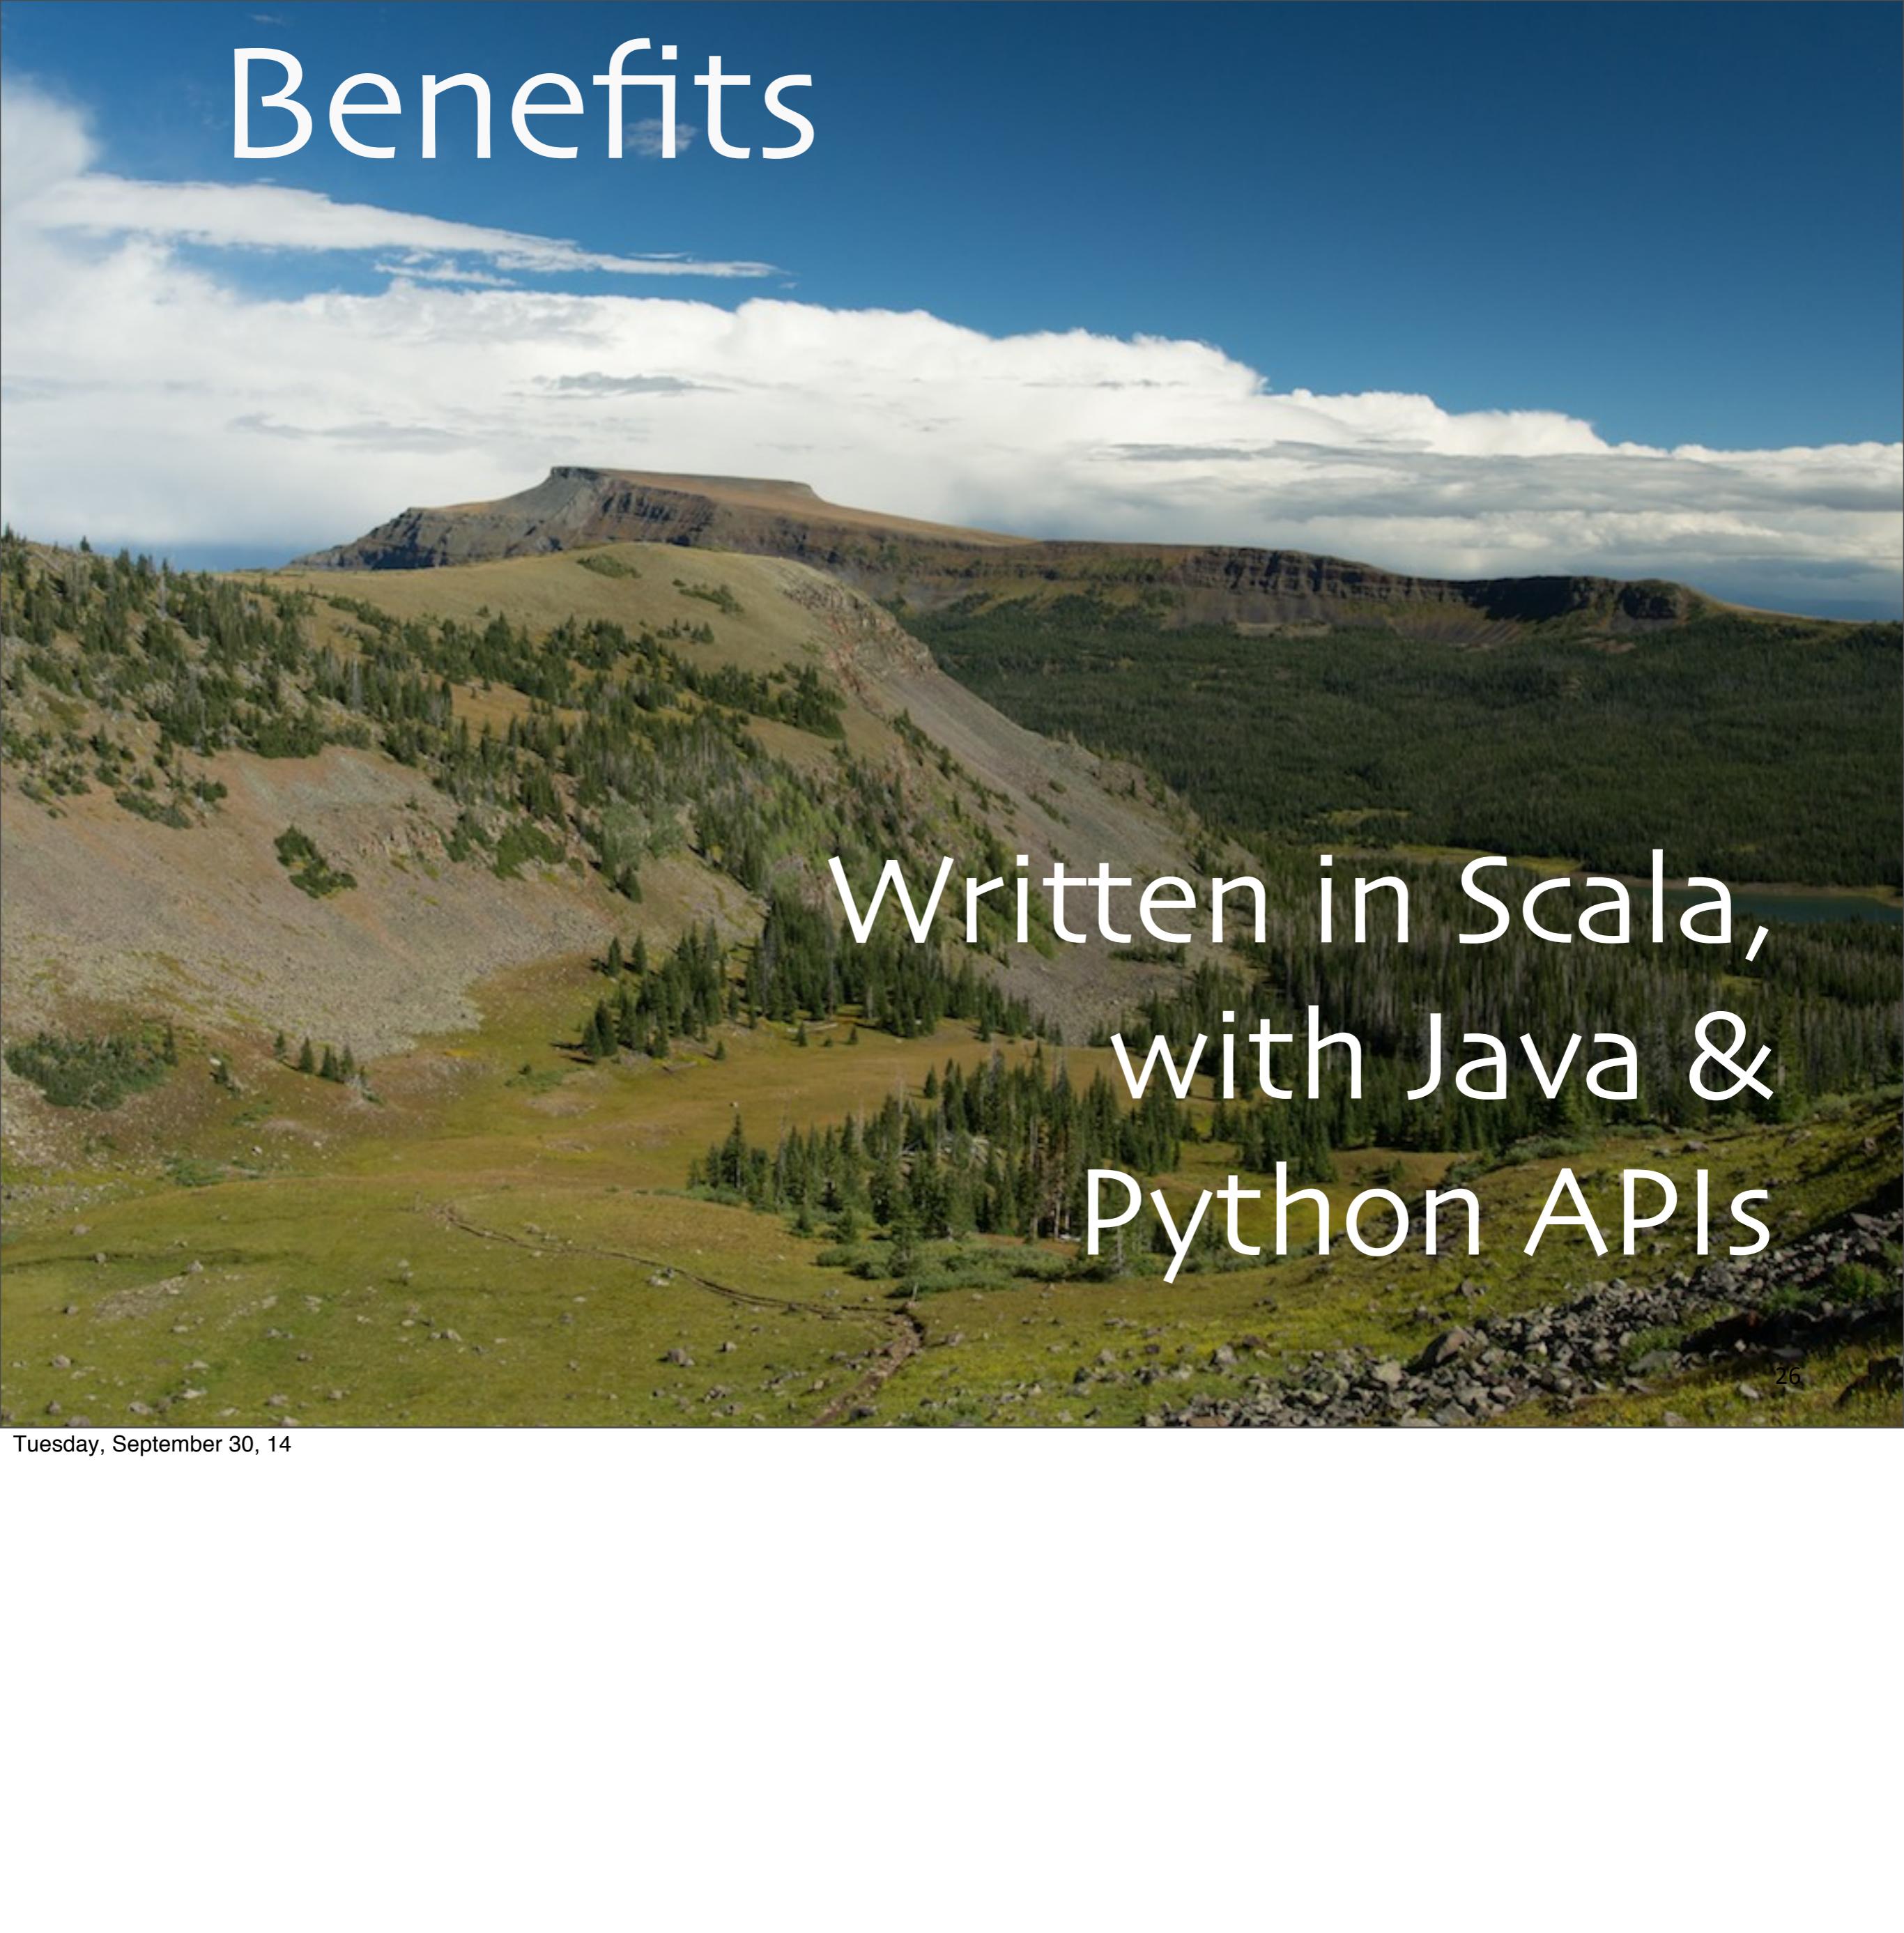
Tuesday, September 30, 14

<http://spark.apache.org>

# Benefits

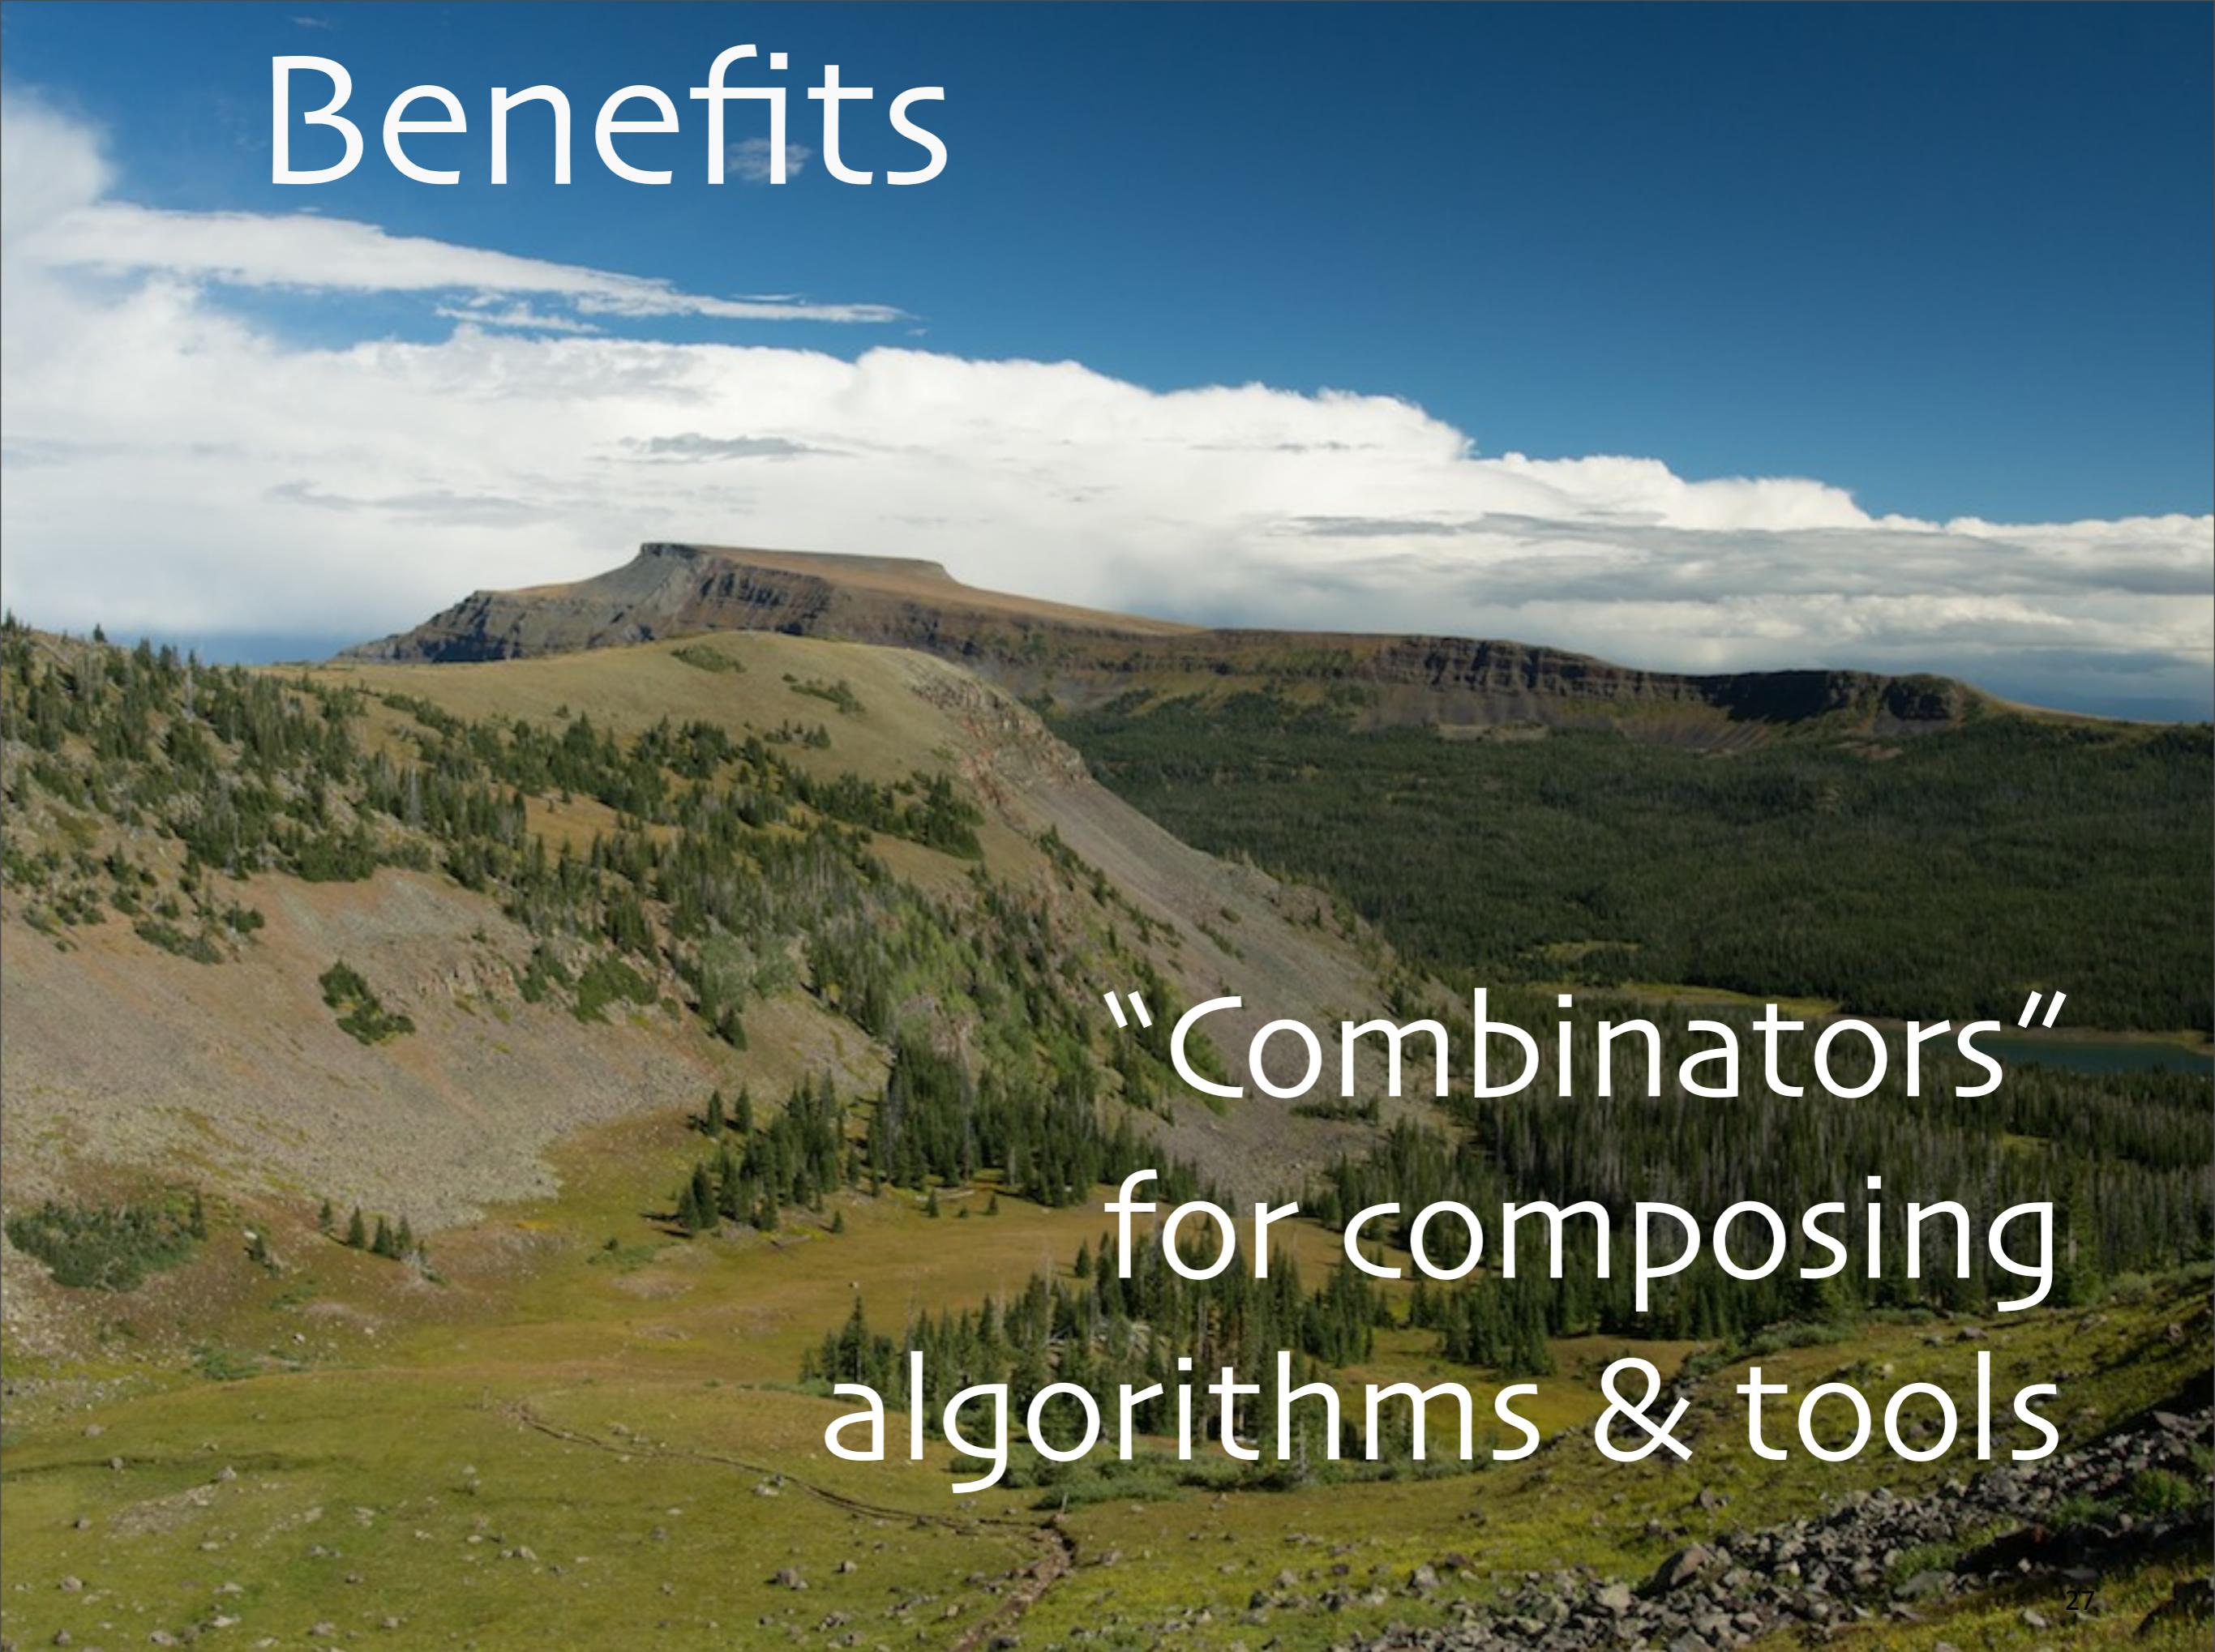
Flexible, elegant, concise  
programming model

# Benefits

A scenic view of a mountain range under a blue sky with white clouds. The mountains have green forests and yellow grassy slopes.

Written in Scala,  
with Java &  
Python APIs

# Benefits



“Combinators”  
for composing  
algorithms & tools

27

Tuesday, September 30, 14

Once you learn the core set of primitives, it's easy to compose non-trivial algorithms with little code.

# Benefits

Many deployment options

Hadoop (YARN)  
Mesos  
EC2  
Standalone

28

Tuesday, September 30, 14

Not restricted to Hadoop, when you don't need it, e.g., because you want to "enhance" existing applications with data analytics, running in the same infrastructure.

# Resilient Distributed Datasets

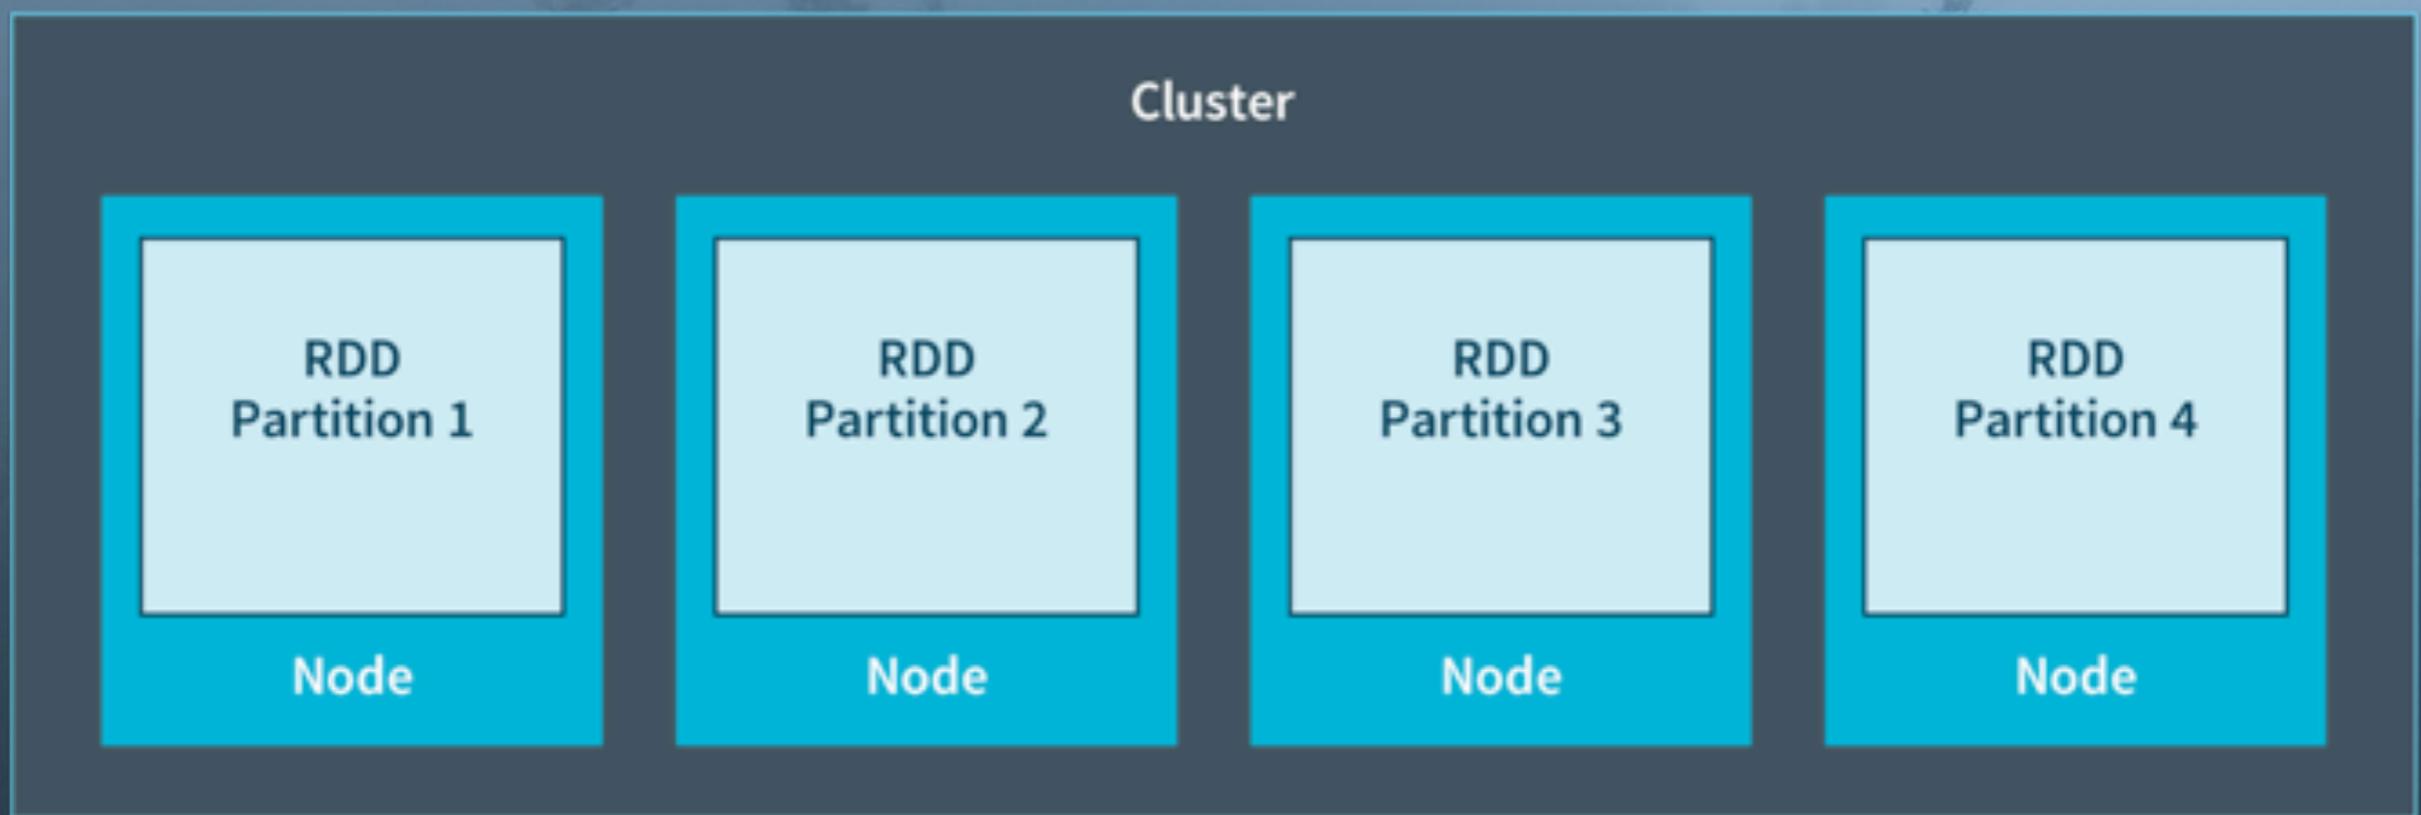


The core abstraction

29

Tuesday, September 30, 14

Data is shared over the cluster in RDDs. This is the core abstraction everyone else builds on.



# Example

Inverted  
Index

31

Tuesday, September 30, 14

Let's see our example rewritten in Spark.

```

import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
          text.split("\\W+").map(s =>

```

32

Tuesday, September 30, 14

It starts with imports, then declares a singleton object (a first-class concept in Scala), with a main routine (as in Java). The methods are colored yellow again. Note this time how dense with meaning they are this time.

You begin the workflow by declaring a SparkContext (in “local” mode, in this case). The rest of the program is a sequence of function calls, analogous to “pipes” we connect together to perform the data flow.

Next we read one or more text files. If “data/crawl” has 1 or more Hadoop-style “part-NNNNN” files, Spark will process all of them (in parallel if running a distributed configuration; they will be processed synchronously in local mode).

sc.textFile returns an RDD with a string for each line of input text. So, the first thing we do is map over these strings to extract the original document id (i.e., file name), followed by the text in the document, all on one line. We assume tab is the separator. “(array(0), array(1))” returns a two-element “tuple”. Think of the output RDD having a schema “String fileName, String text”.

```

}.flatMap {
    case (path, text) =>
        text.split("""\W+""") map {
            word => (word, path)
        }
}
.map {
    case (w, p) => ((w, p), 1)
}
.reduceByKey {
    case (n1, n2) => n1 + n2
}
.map {
    case ((w, p), n) => (w, (p, n))
}
.groupBy {
    case (w, (p, n)) => w
}

```

33

Tuesday, September 30, 14

flatMap maps over each of these 2-element tuples. We split the text into words on non-alphanumeric characters, then output collections of word (our ultimate, final “key”) and the path. Each line is converted to a collection of (word,path) pairs, so flatMap converts the collection of collections into one long “flat” collection of (word,path) pairs.

Then we map over these pairs and add a single count of 1.

reduceByKey does an implicit “group by” to bring together all occurrences of the same (word, path) and then sums up their counts. Note the input to the next map is now ((word, path), n), where n is now  $\geq 1$ . We transform these tuples into the form we actually want, (word, (path, n)).

```

}
  .groupBy {
    case (w, (p, n)) => w
  }
  .map {
    case (w, seq) =>
      val seq2 = seq map {
        case (_, (p, n)) => (p, n)
      }
      (w, seq2.mkString(", "))
  }
  .saveAsTextFile(argz.outpath)

sc.stop()
}
}

```

34

Tuesday, September 30, 14

Now we do an explicit group by to bring all the same words together. The output will be (word, (word, (path1, n1)), (word, (path2, n2)), ...). The last map removes the redundant “word” values in the sequences of the previous output. It outputs the sequence as a final string of comma-separated (path,n) pairs.

We finish by saving the output as text file(s) and stopping the workflow.

```
}

.map {
  case (w, p) => ((w, p), 1)
}

.reduceByKey {
  case (n1, n2) => n1 + n2
}

.map {
  case ((w, p), n) => (w, (p, n))
}

.groupBy {
  case (w, (p, n)) => w
}

.map {
  case (w, seq) =>
    val seq2 = seq map {
      case (_, (p, n)) => (p, n)
```

Chain  
combinators  
together

$$\nabla \cdot \mathbf{D} = \rho$$

$$\nabla \cdot \mathbf{B} = 0$$

$$\nabla \times \mathbf{E} = -\frac{\partial \mathbf{B}}{\partial t}$$

$$\nabla \times \mathbf{H} = \mathbf{J} + \frac{\partial \mathbf{D}}{\partial t}$$

36

Tuesday, September 30, 14

Another example of a beautiful and profound DSL, in this case from the world of Physics: Maxwell's equations: <http://upload.wikimedia.org/wikipedia/commons/c/c4/Maxwell'sEquations.svg>

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
          text.split("""\W+""") map {
            word => (word, path)
          }
      }
      .map {
        case (w, p) => ((w, p), 1)
      }
      .reduceByKey {
        case (n1, n2) => n1 + n2
      }
      .map {
        case ((w, p), n) => (w, (p, n))
      }
      .groupByKey {
        case (w, (p, n)) => w
      }
      .map {
        case (w, seq) =>
          val seq2 = seq map {
            case (_, (p, n)) => (p, n)
          }
          (w, seq2.mkString(", "))
      }
      .saveAsTextFile(argz.outpath)

    sc.stop()
  }
}
```

# Altogether



That version took me  
~30 mins. to write

38

Tuesday, September 30, 14

When you have a concise, flexible API, you can turn a “software development project” into a script! It transforms your productivity.

# RDDs + Core APIs:

A foundation for  
other tools

39

Tuesday, September 30, 14

The good API also provides an excellent foundation for other tools to build on.

# Extensions

MLlib  
GraphX  
Tachyon

...

40

Tuesday, September 30, 14

MLlib - a growing library of machine learning algorithms.

GraphX - for representing data as a graph and applying graph algorithms to it.

Tachyon - an experiment to generalize Spark's caching mechanism into a standalone service, so data is shareable between apps and more durable. I believe it will be transformative!

# Extensions

..  
Spark SQL

...

41

Tuesday, September 30, 14

Let's look at the SQL abstractions layered on top.



# RDD API + SQL

42

Tuesday, September 30, 14

Best of both worlds: SQL for concision, the RDD API for Turing-complete, general-purpose computing. Also adds elegant handling of the “schema” for data.



# Hive Interop

43

Tuesday, September 30, 14

Let's us query Hadoop Hive "tables". We can create or delete them, too.



# JSON

44

Tuesday, September 30, 14

New feature. Can read JSON records and infer their schema. Can write RDD records as JSON.

# Example

Use the Crawl data  
for Inverted Index

45

Tuesday, September 30, 14

It's a bit tricky to use the inverted index data, because of the variable list of (docid, N) values, so we'll use the crawl data, which is easier for our purposes.

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
import org.apache.spark.sql.{  
    SQLContext, SchemaRDD}
import org.apache.spark.rdd.RDD

case class CrawlRecord(  
    docid: String, contents: String)

def makeCrawlRecord(line: String) = {...}

def dosql(qry: String, n: Int = 100) =  
    sql(qry).collect.take(n) foreach println

val crawlData = "/path/to/directory"

val sc = new SparkContext("...","Crawl")
```

46

Tuesday, September 30, 14

Starts out like a typical Spark program...

Defines a “case class” (think normal Java class where the args are automatically turned into fields) to represent each record.

Defines a method to parse each input line and output a CrawlRecord (details omitted)

Defines a helper method to take a SQL query as a string, run it using the “sql(...)” method provided by SparkSQL, grab the first n elements and print them, one per line.

Finally, defines the path to the input for the crawl data.

```
val sc = new SparkContext("...", "Crawl")
```

```
val crawl = for {
    line <- sc.textFile(crawlData)
    cr <- makeCrawlRecord(line)
} yield cr
```

```
crawl.registerAsTable("crawl")
crawl.printSchema
```

```
dosql("""
    SELECT docid, contents FROM crawl
    LIMIT 10""")
```

```
val crawlPerWord = crawl flatMap {
    case CrawlRecord(docid, contents) =>
        contents.trim.split("""[^'\w']""") map
```

47

Tuesday, September 30, 14

The for loop takes each line of input and parses it into a CrawlRecord. So, “crawl” has the type RDD[CrawlRecord]. Then we register it as a table and print its schema.

Now run a query!

```
val crawlPerWord = crawl flatMap {  
  case CrawlRecord(docid, contents) =>  
    contents.trim.split("""[^\\w']""") map  
(word => CrawlRecord(docid, word))  
}
```

```
crawlPerWord.registerAsTable("crawl2")  
crawlPerWord.cache
```

```
dosql("SELECT * FROM crawl2 LIMIT 10")  
dosql("""  
  SELECT DISTINCT * FROM crawl2  
  WHERE contents = 'management'""")  
dosql("""  
  SELECT contents, COUNT(*) AS c  
  FROM crawl2  
  GROUP BY contents
```

48

Tuesday, September 30, 14

Create a new, similar RDD where the records are the docids with each word in the document, not one record for each docid with the entire contents.

Register this RDD as a table and this time cache it in memory, since we'll write several queries against it.

```
crawlPerWord.registerAsTable("crawl2")
crawlPerWord.cache
```

```
dosql("SELECT * FROM crawl2 LIMIT 10")
dosql("""
    SELECT DISTINCT * FROM crawl2
    WHERE contents = 'management'""")
dosql("""
    SELECT contents, COUNT(*) AS c
    FROM crawl2
    GROUP BY contents
    ORDER BY c DESC LIMIT 100""")
```

# Extensions

... and  
Streaming.

# Capture & process event time slices

Tuesday, September 30, 14

A clever extension to the existing batch-oriented RDD model; use smaller batches! So, it's not a replacement for true event-processing systems, like Storm, message queues.

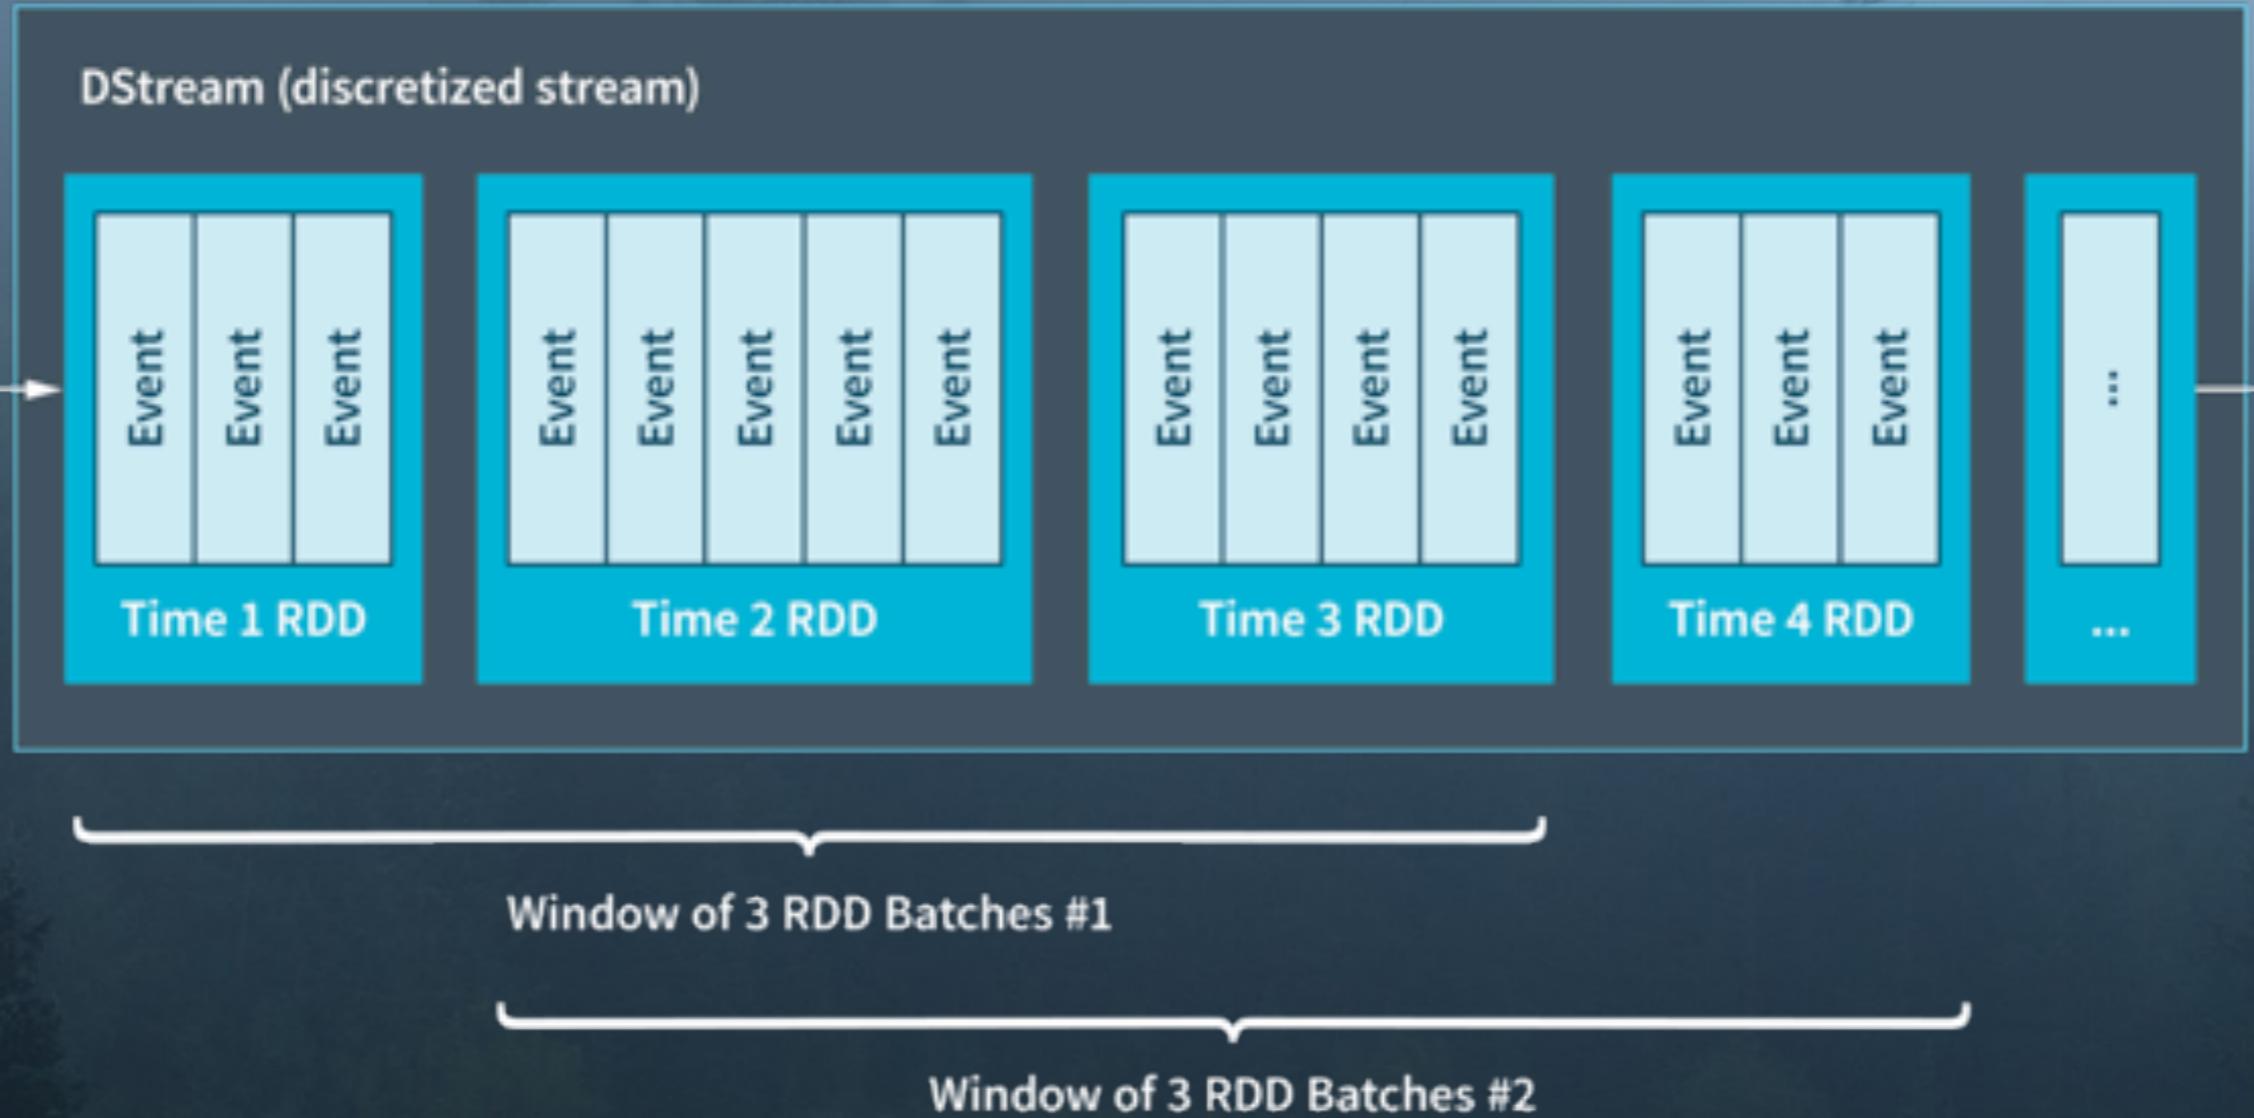


Each slice is  
an RDD.  
Plus window  
functions



Tuesday, September 30, 14

We get all the familiar RDD functions, “for free”, plus functions for working with windows of batches.



# Example

Use “live” Crawl data  
for Inverted Index

54

Tuesday, September 30, 14

Continue using the crawl data, but “pretend” we’re reading it live from a socket.

```
// ... imports, etc.  
val sc = new SparkContext(...)  
val ssc = new StreamingContext(  
    sc, Seconds(60)) ← "Batch" size  
ssc2.addStreamingListener(  
    /*... listener for end of data ...*/)  
val sqlc = new SQLContext(sc)  
import sqlc._
```

```
val inputDStream =  
    sc.socketTextStream(server,  
port).flatMap(_.split("\n"))
```

```
val crawlWords = for {  
    line <- inputDStream  
    cr1 <- makeCrawlRecord(line)
```

55

Tuesday, September 30, 14

We won't show everything now, just the interesting bits...

We create a SparkContext, then wrap it with a new StreamingContext object, where we'll grab the records in 60-second increments, AND a SQLContext as before (optional).

We also add listener for stream events, such as end of input (details omitted).

Finally, create a DStream that will ingest events from a socket of plain text data. It uses flatMap to split the input stream into "records" on newlines.

```
port).flatMap(_.split("\n"))

val crawlWords = for {
    line <- inputStream
    cr1 <- makeCrawlRecord(line)
    word <- cr1.contents.trim.split(
        """[^w']""")
} yield (CrawlRecord(cr1.docid, word))

crawlWords.window(
    Seconds(300), Seconds(60))
.foreachRDD { rdd =>
    rdd.registerAsTable("crawlWords")
    dosql("""
        SELECT contents, COUNT(*) AS c
        FROM crawlWords
        GROUP BY contents
    """)
```

56

Tuesday, September 30, 14

As before, parse each line into a CrawlRecord of docids and contents, then immediately convert to docids and individual words. (We used separate RDDs for these two steps previously.)

```
} yield (CrawlRecord(cr1.docid, word))
```

```
crawlWords.window(  
    Seconds(300), Seconds(60))  
.foreachRDD { rdd =>  
    rdd.registerAsTable("crawlWords")  
    dosql("""  
        SELECT contents, COUNT(*) AS c  
        FROM crawlWords  
        GROUP BY contents  
        ORDER BY c DESC LIMIT 100""")  
}
```

Window size,  
"skip size.

```
ssc2.start()  
ssc2.awaitTermination()
```

57

Tuesday, September 30, 14

Now process event batches in windows of 300 seconds (5 batches of 60 seconds) and each time we move the window, move by 60 seconds (one batch). Then, for each RDD, register as a table and run a query.

Finally, start processing and wait for it to finish (if ever!).

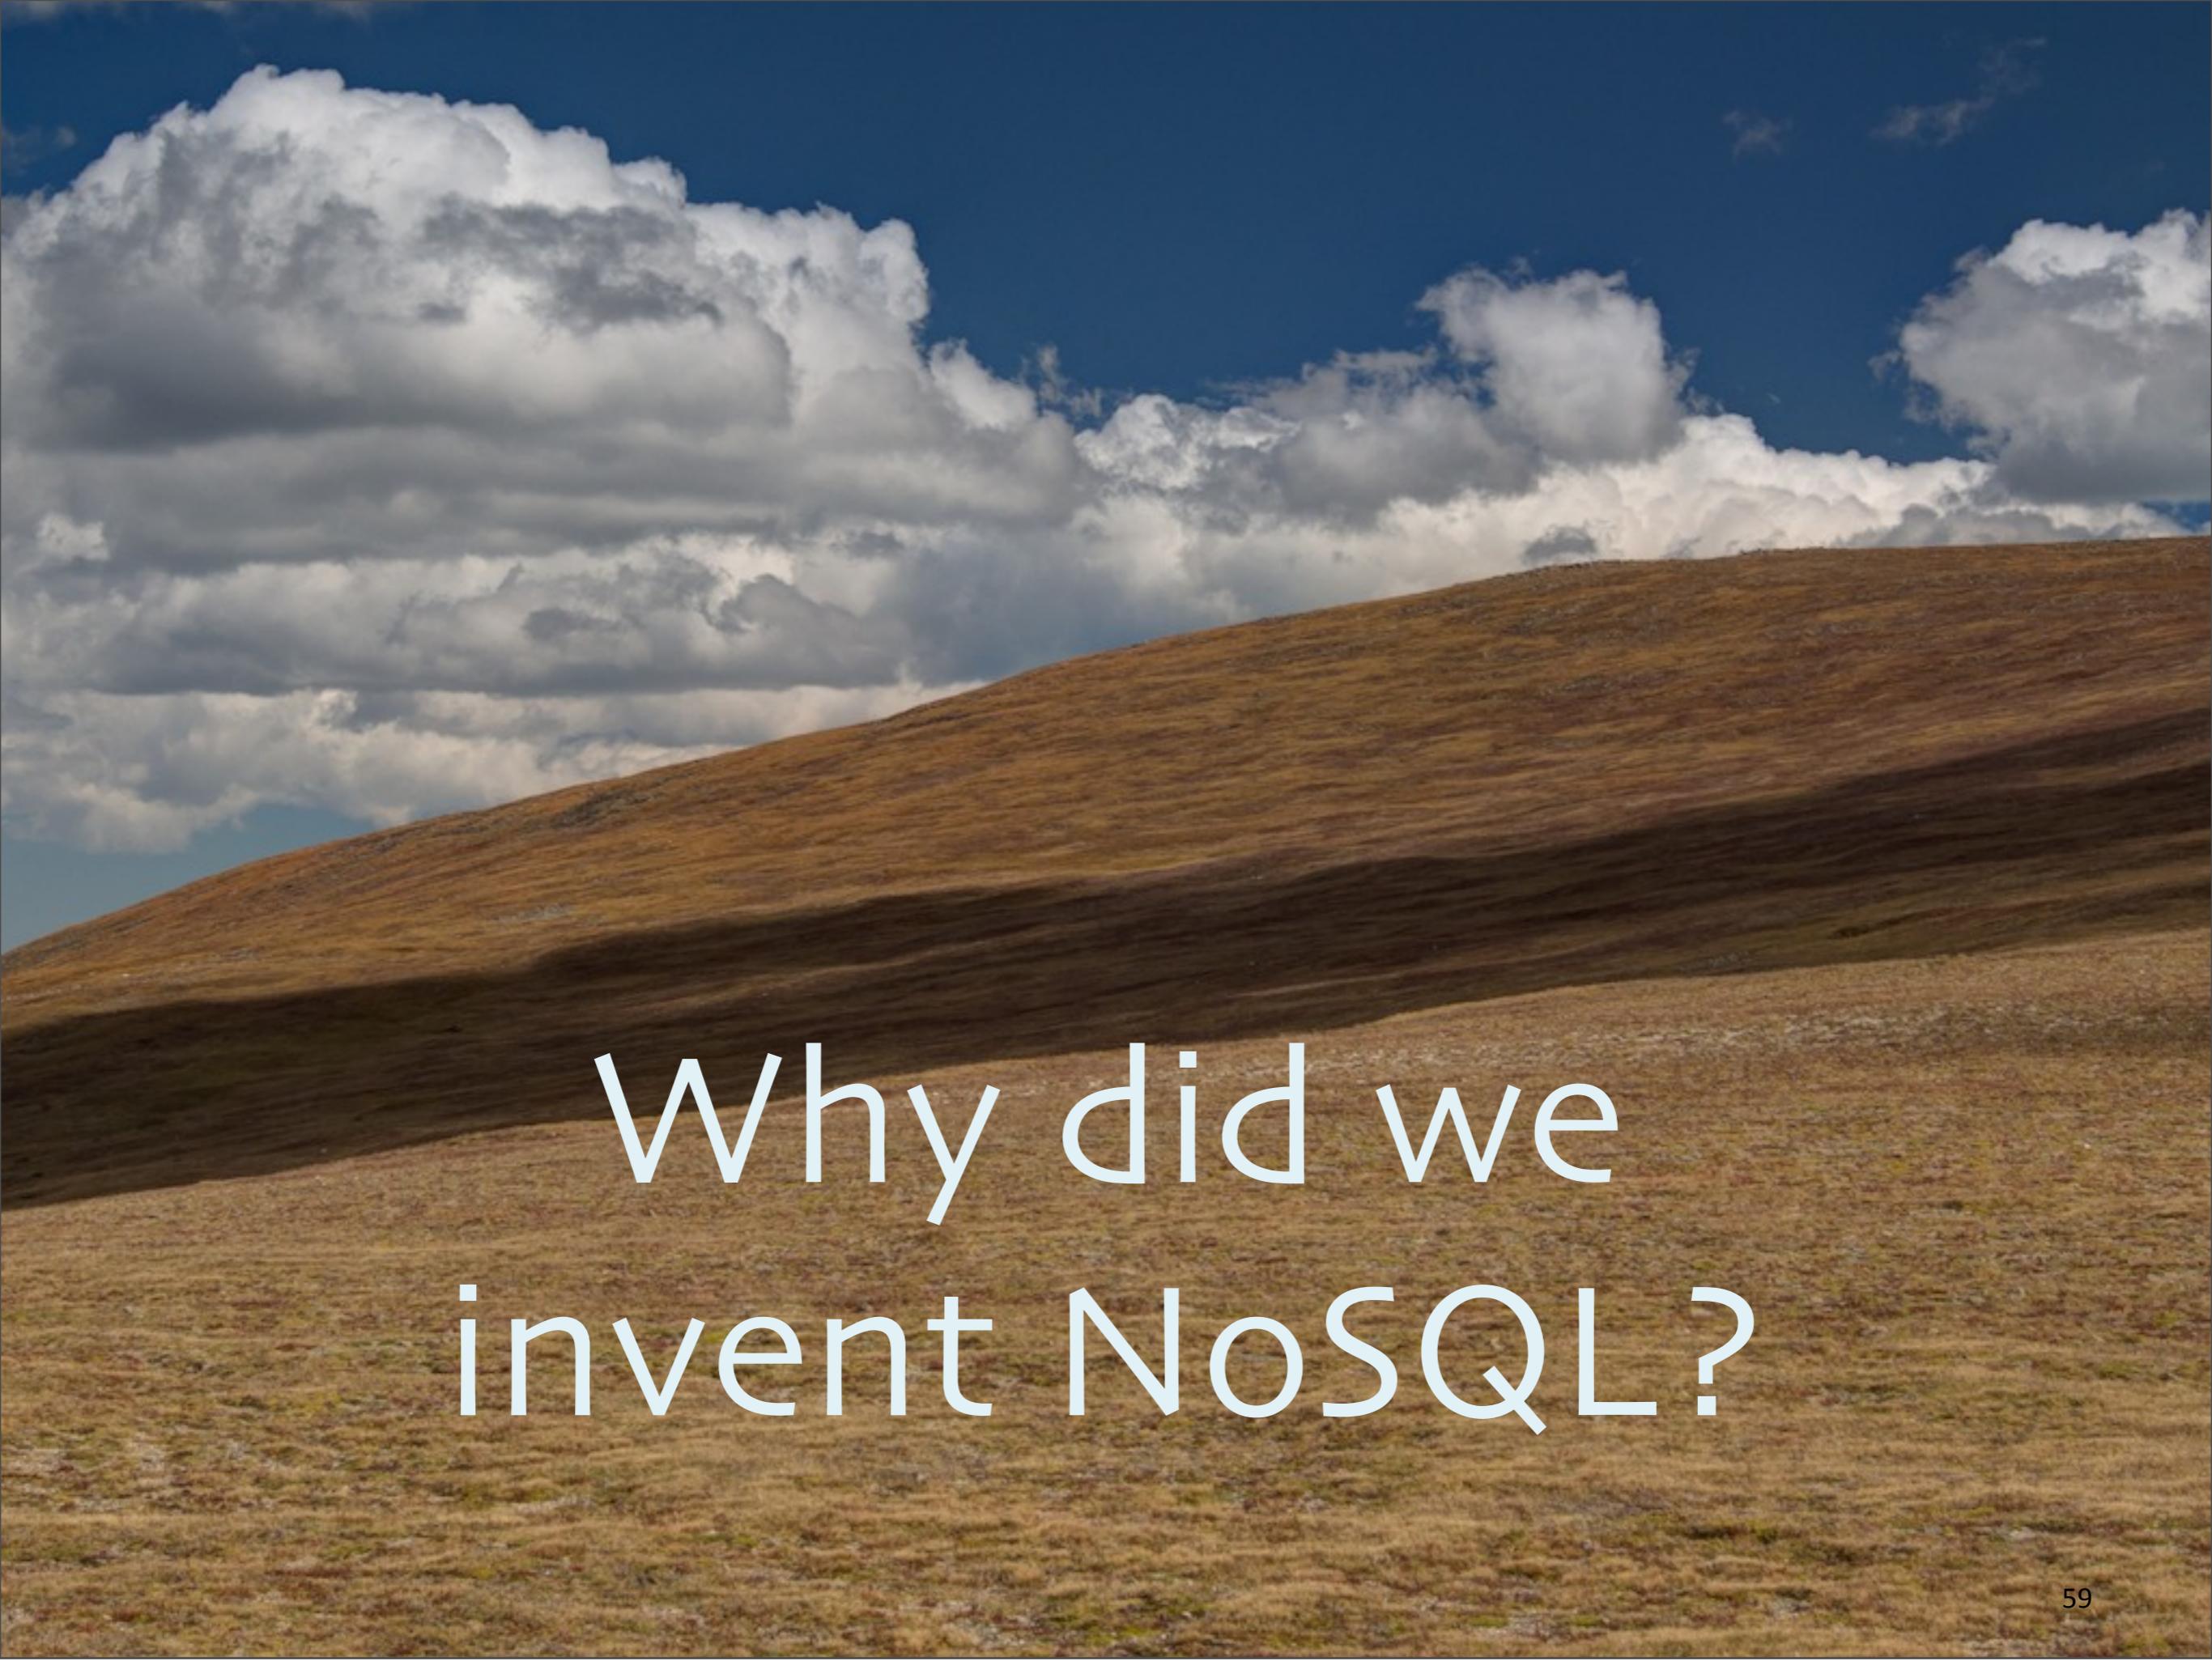
# Return of SQL!



58

Tuesday, September 30, 14

So, SQL is very useful for “structured” data in Hadoop. In fact, SQL has experienced a renaissance in Big Data



# Why did we invent NoSQL?

59

Tuesday, September 30, 14

First, why did NoSQL emerge in the first place?

# why NoSQL?

Massive Scale

60

Tuesday, September 30, 14

1. We needed to manage unprecedented data set sizes, economically. Existing Relational tools couldn't handle the size, especially at low cost.

# why NoSQL?

CAP

61

Tuesday, September 30, 14

Sometimes remaining available and accepting eventual consistency is the tradeoff we want when partitions occur. Relational is CP, if a partition happens we prefer consistency, so the DB won't be available until the partition is resolved. But many apps can accept lack of consistency if they can still remain available.

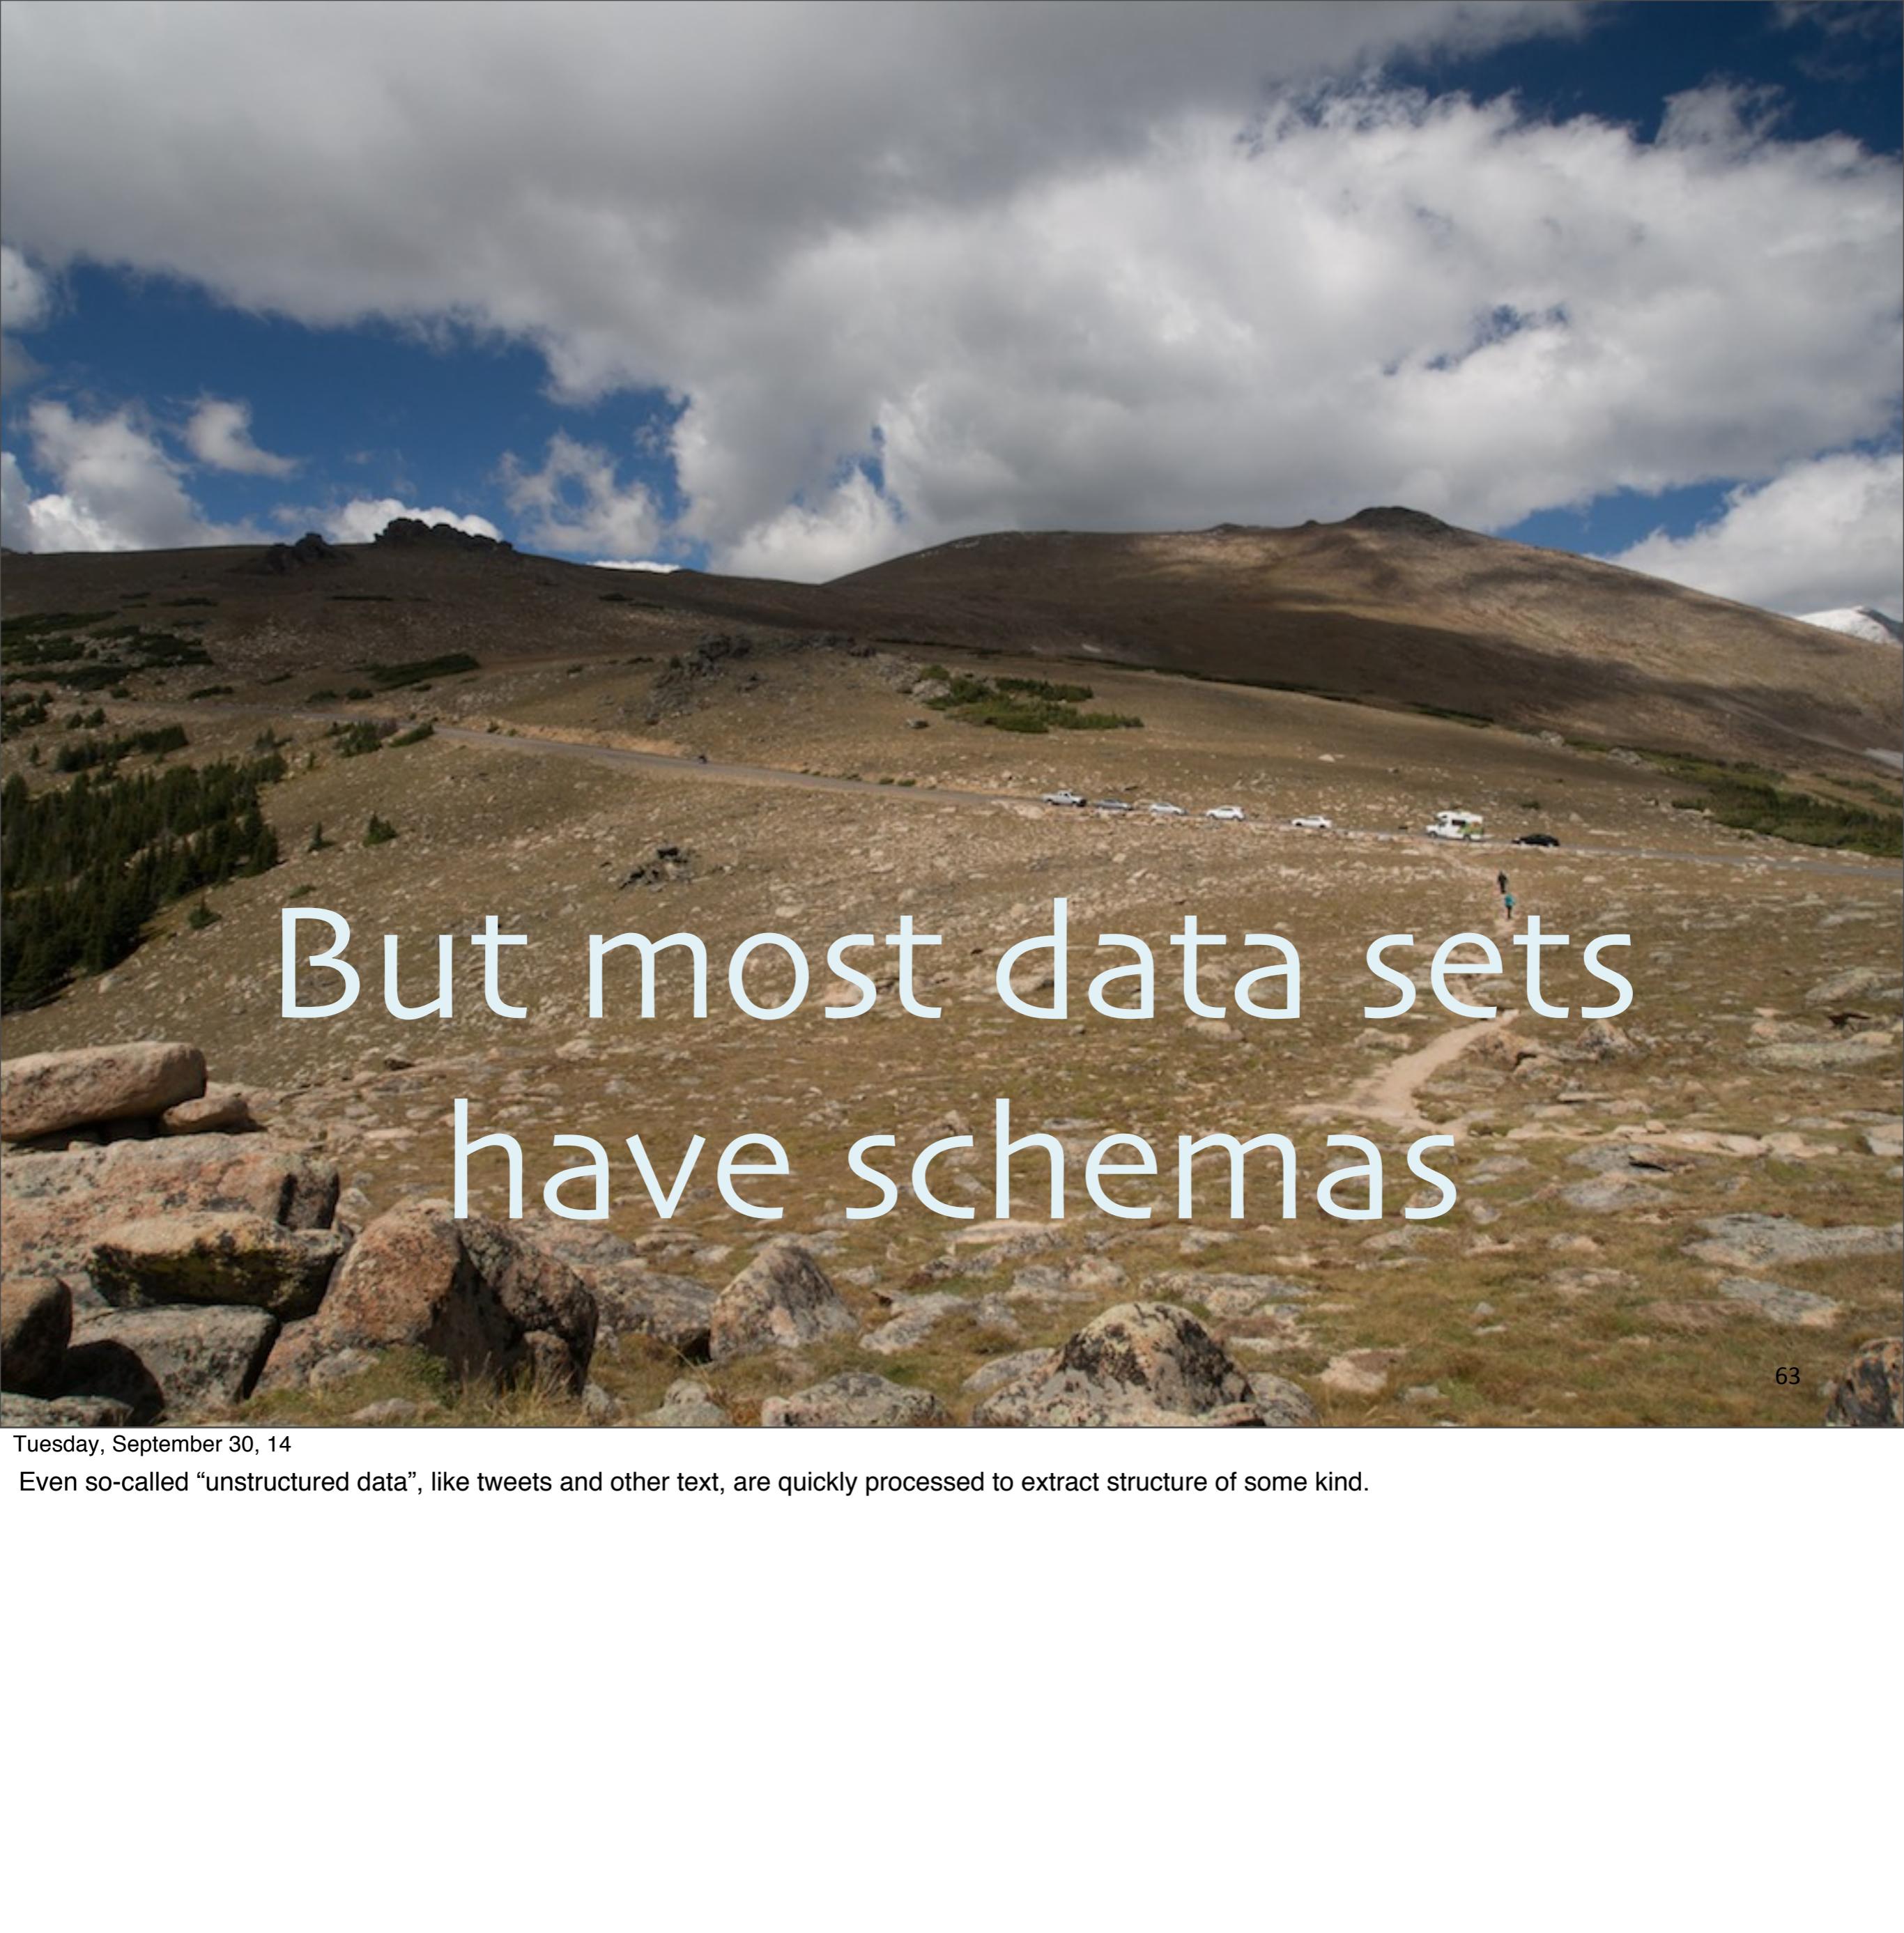
# Why NoSQL?

Not all data is  
relational

62

Tuesday, September 30, 14

Key-value stores, hierarchical data, e.g.,JSON/XML docs, etc. are alternative forms that work better for many scenarios.

A scenic view of a mountain landscape under a blue sky with white clouds. In the foreground, there are large rocks and patches of green and brown grass. A dirt road winds its way through the terrain, leading towards a cluster of vehicles and people in the middle ground. The background features rolling hills and mountains.

But most data sets  
have schemas

63

Tuesday, September 30, 14

Even so-called “unstructured data”, like tweets and other text, are quickly processed to extract structure of some kind.

# Two New APProaches for SQL

64

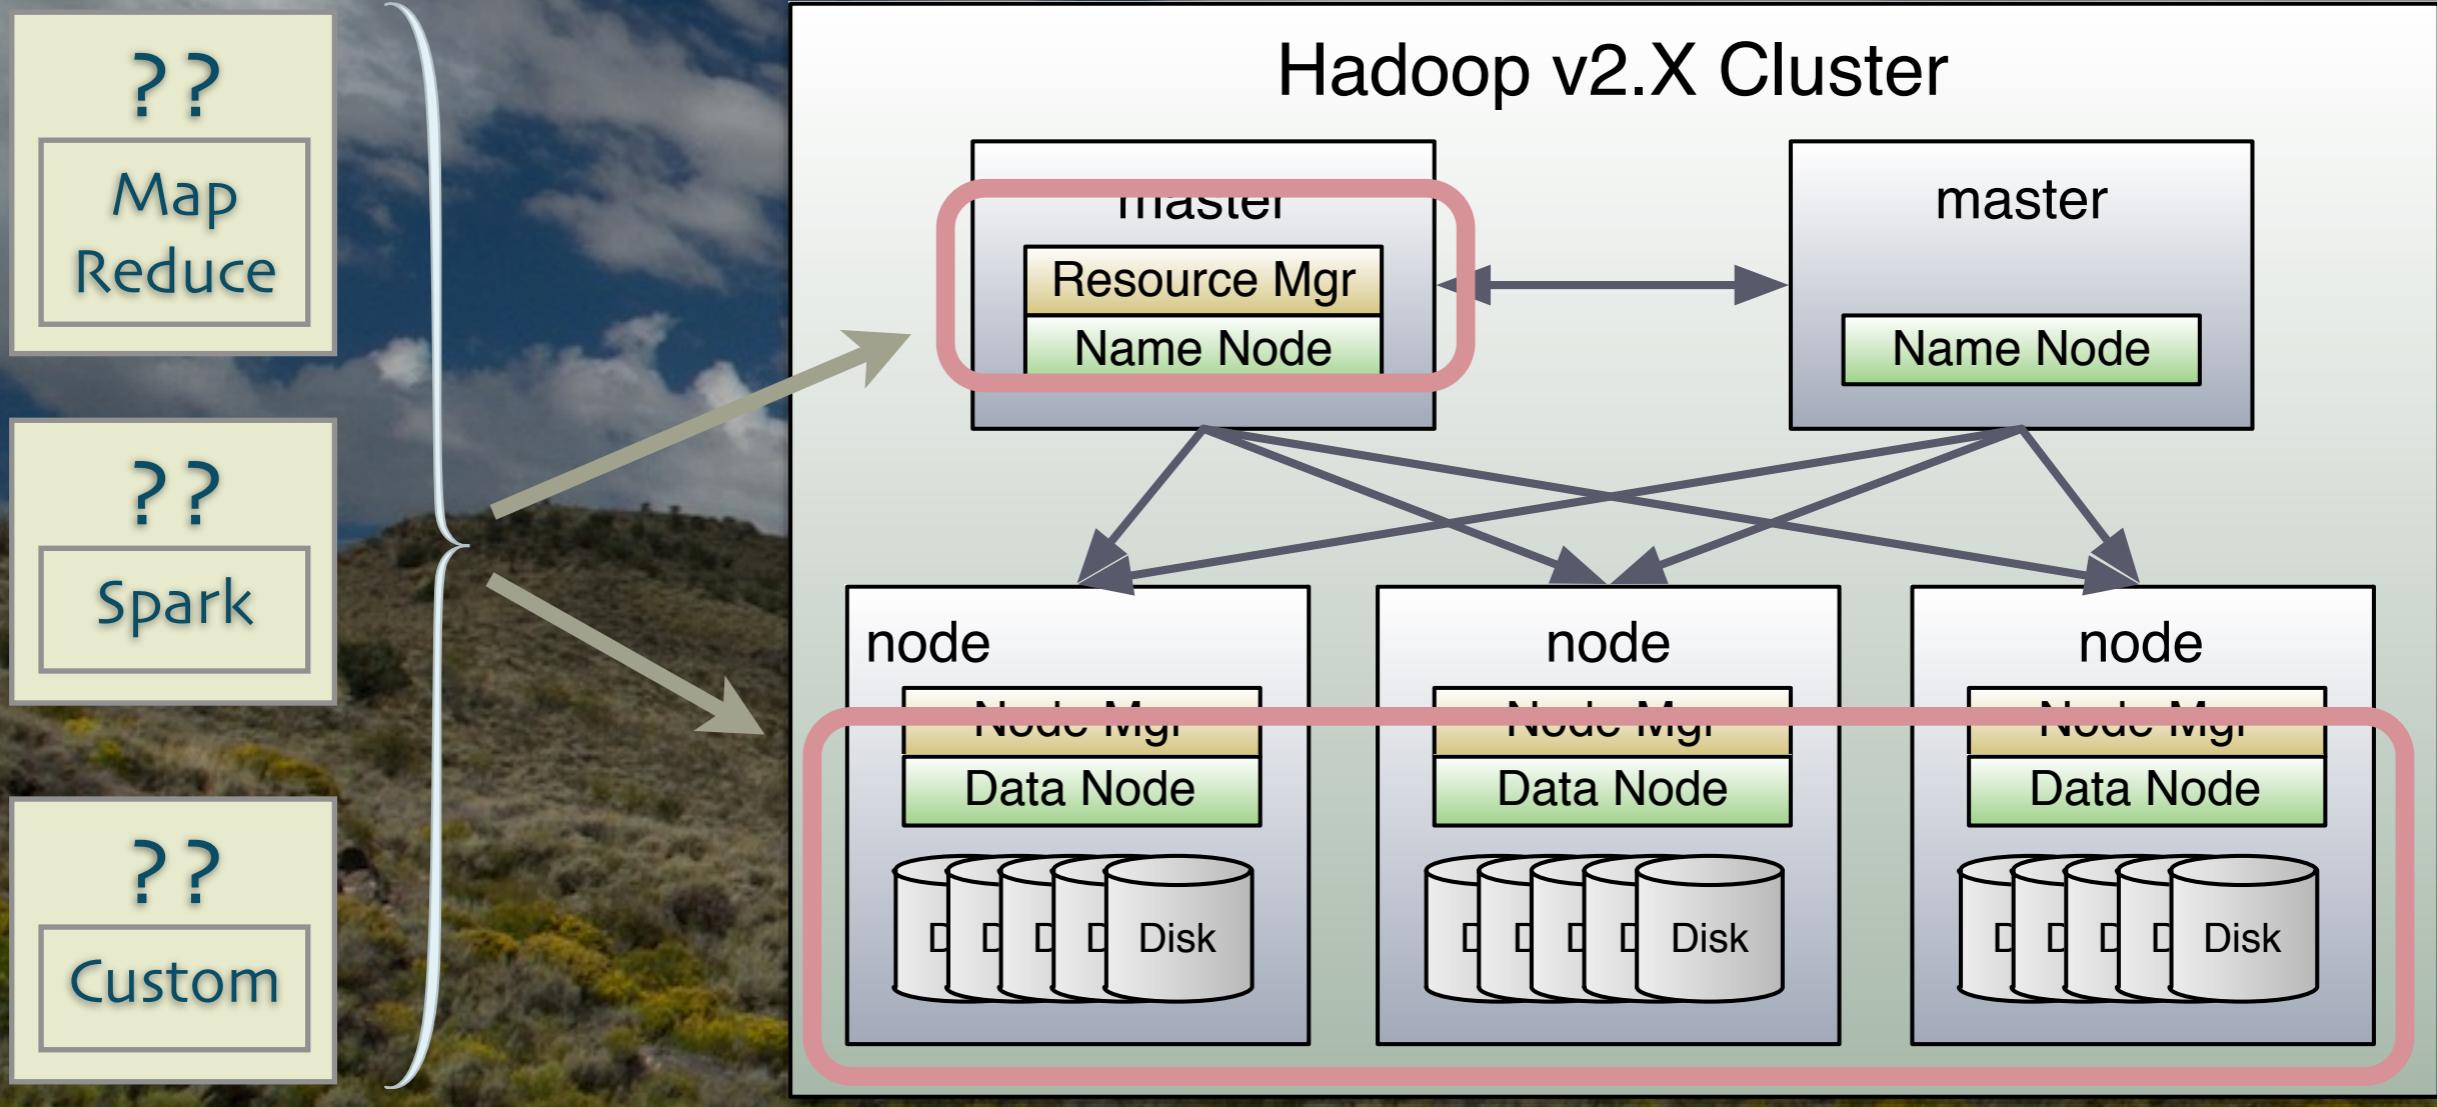
# 1. Query Engine + HDFS

65

Tuesday, September 30, 14

First idea, put SQL-based query abstractions on top of simple storage, like flat files in HDFS, MapRFS, S3, etc. The query abstractions can be a “DSL” implemented in a generic framework like MapReduce or Spark, or with a custom query engine optimized for the job.

# 1. Query Engine + HDFS

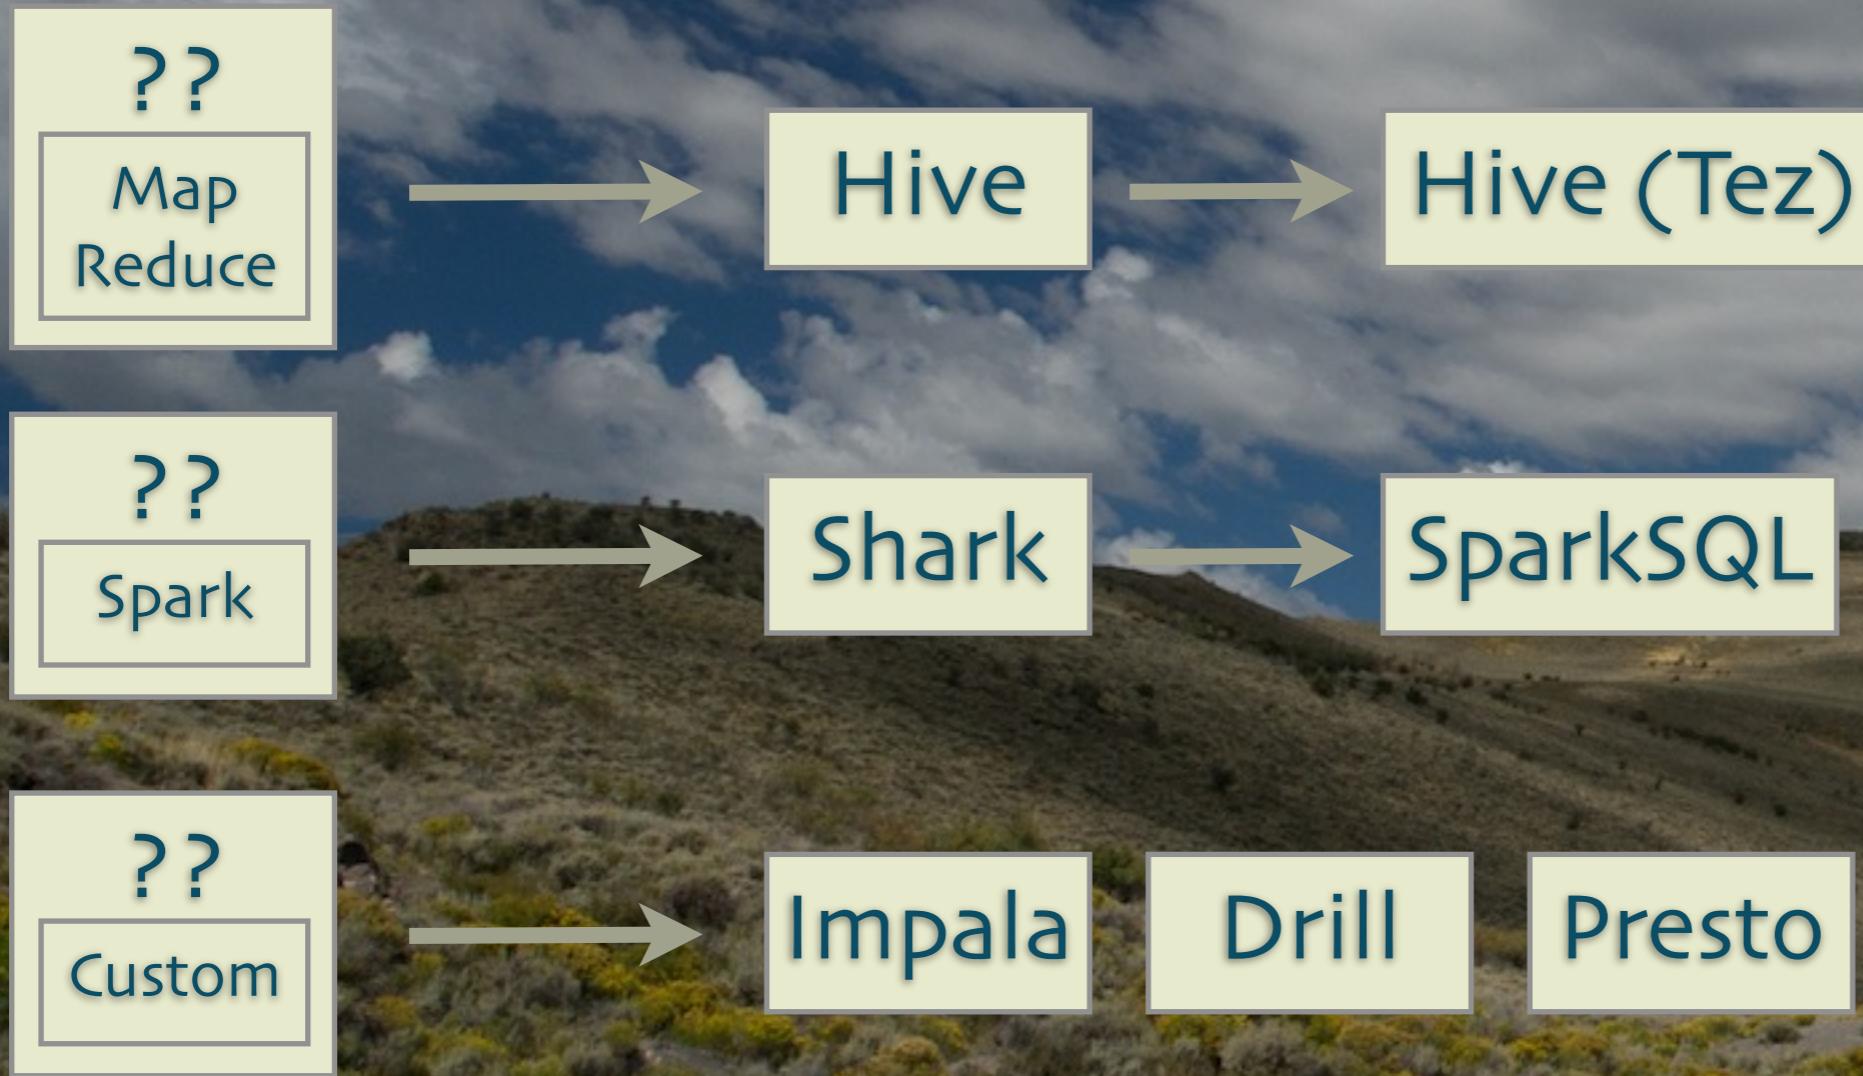


66

Tuesday, September 30, 14

You could write engines in MR, Spark, or something custom, which may offer less flexibility, but better performance.

# 1. Query Engine + HDFS



67

Tuesday, September 30, 14

Hive, developed at Facebook, pioneered SQL on Hadoop. It has recently been ported to a new, higher-performance engine called Tez. Tez is a competitor to Spark, but isn't gaining the same traction.

The Spark team ported Hive to Spark and achieved 30x+ performance improvements. Shark is now deprecated; it's being replaced with a better engineered query engine called Catalyst, inside SparkSQL.

Impala was the first custom query engine, inspired by Google's Dremel. It holds the current performance records for SQL on Hadoop. Presto is a Facebook project. Drill is another Apache project. It's more portable across Hadoop platforms than Impala, but not as fast.



## 2. NewSQL

68

Tuesday, September 30, 14

These are new, relational databases, separate from Hadoop altogether. They leverage the scalability and resiliency lessons of NoSQL, but restore the relational model.



Google Spanner, F1  
VoltDB  
NuoDB  
FoundationDB  
SAP HANA

## 2. NewSQL

69

Tuesday, September 30, 14

Google Spanner is globally distributed with \*global transactions\*. The others are commercial projects.

# Looking Ahead



70

Tuesday, September 30, 14

Gazing into the foggy future...

# Spark + Mesos



71

Tuesday, September 30, 14

Mesos may be all that many teams need, if they don't already have Hadoop (YARN), and especially if they have other infrastructure running on Mesos.  
(Technically, you can run Hadoop on Mesos, too.)

# Flexible cloud deployments

72

Tuesday, September 30, 14

In general, people are pursuing flexible ways of deploying big data apps, especially when they integrate with other systems running in different cloud or virtualization environments.



# Watch Tachyon

73

Tuesday, September 30, 14

I think Tachyon will be disruptive when it's mature.

# H<sub>2</sub>O

<https://github.com/0xdata/h2o>

74

Tuesday, September 30, 14

Check out this high-performance computing engine. <https://github.com/0xdata/h2o> They are integrating it with Spark.  
See Cliff Click's talk!

# Recap



75

Tuesday, September 30, 14

# Spark

Replaces  
MapReduce

# Spark

Supports  
Streaming  
and Batch



Spark  
Integrates  
SQL, ML, & Graph  
libraries

# SQL

Works great in  
Hadoop!

# SQL

NewSQL DBs  
improve Relational  
scalability

# SQL

The world needs  
NoSQL and  
NewSQL

# Prediction

Mesos-based  
environments  
will grow.

# More Stuff by Me...

The screenshot shows a web browser window with the following details:

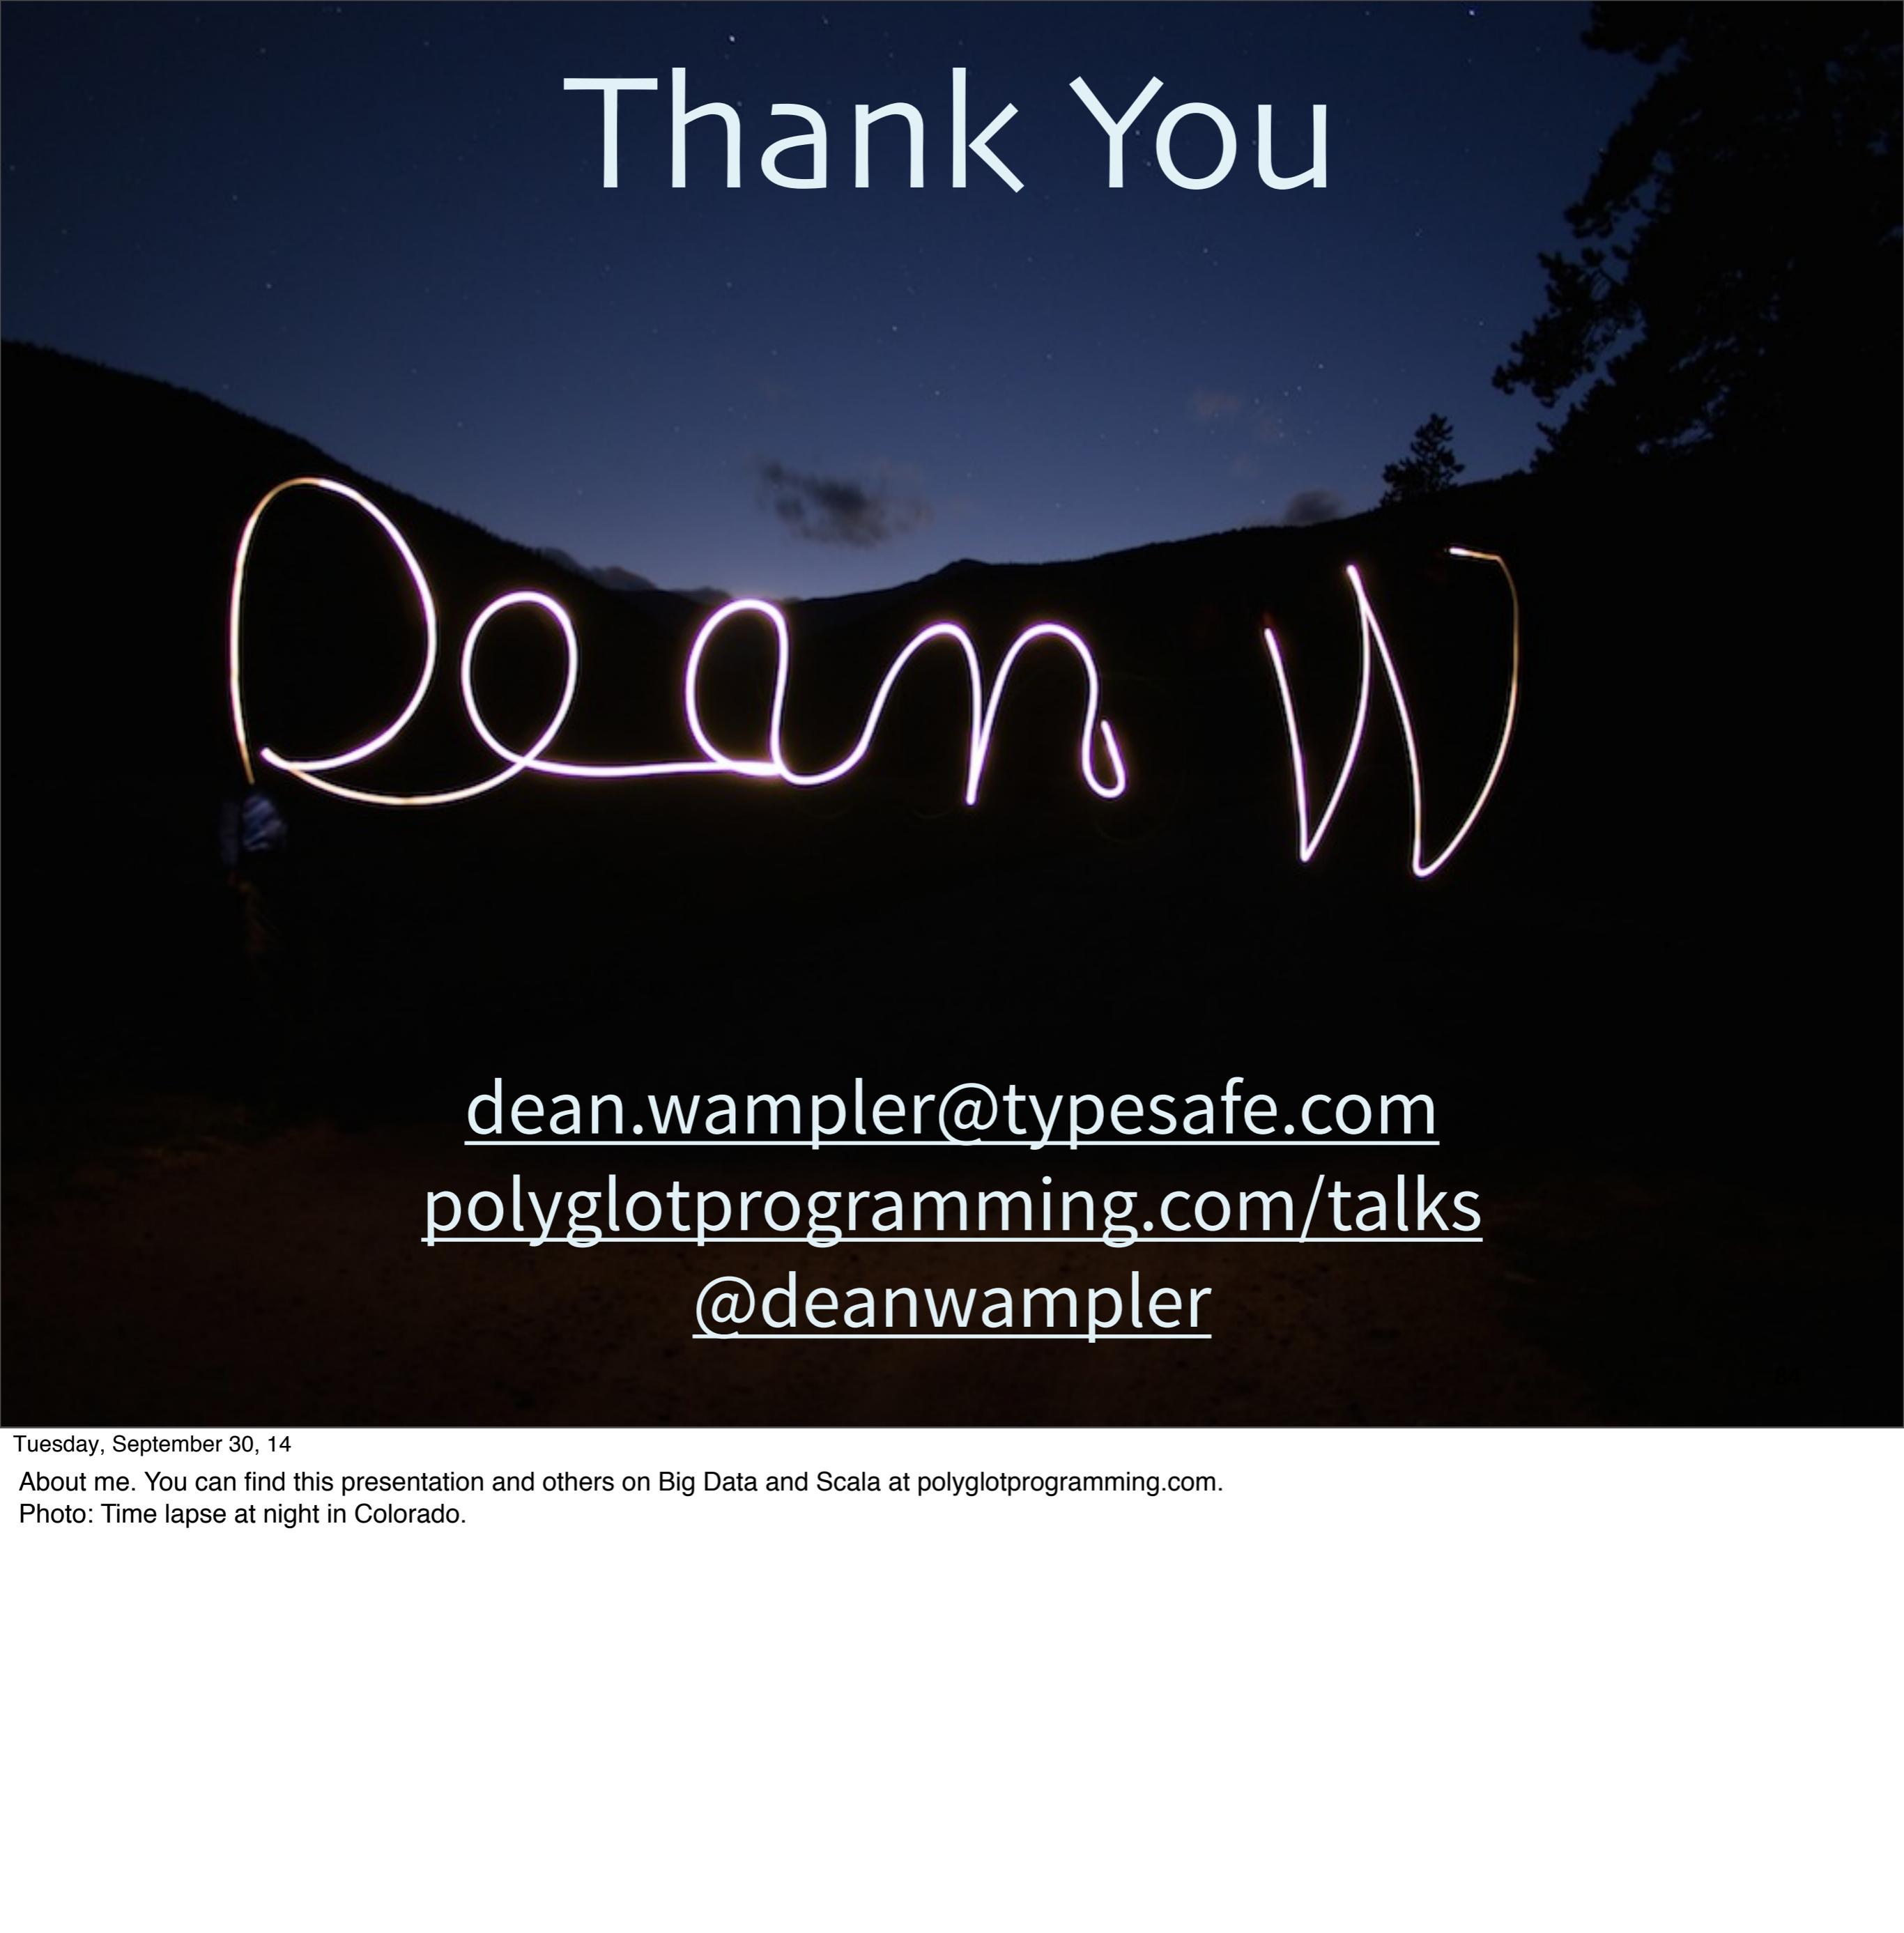
- Address Bar:** Shows the URL [typesafe.com/platform/reactive-big-data/spark](https://typesafe.com/platform/reactive-big-data/spark).
- Page Header:** The page is titled "Reactive Big Data / OVERVIEW".
- Left Sidebar (Red Box):** Labeled "TYPESAFE ACTIVATOR" at the top. It contains the following links:
  - Overview
  - Get Started** (highlighted with a white background)
  - Documentation
  - Browse Templates
  - Contribute Template
- Main Content Area:** The title "Apache Spark and the Typesafe Reactive Platform" is displayed. Below it, a paragraph reads:

At Typesafe, we're committed to helping developers build massively scalable applications on the JVM. Because our users are increasingly building **Reactive** applications that leverage Big Data, we decided to team with Databricks to help developers understand how to better utilize Spark.
- Footer:** A small section at the bottom left includes the date "Tuesday, September 30, 14" and a note about a "1-day Spark Workshop" teaching for Typesafe.

Tuesday, September 30, 14

I have a 1-day Spark Workshop I'm teaching for Typesafe. See this page for details, as well as a whitepaper and blog post on Spark that I wrote.

# Thank You

A photograph of a night sky with mountains in the background. Several bright, glowing lines form the letters "Dean W" across the center of the image, suggesting light painting or a time-lapse effect.

Dean W

[dean.wampler@typesafe.com](mailto:dean.wampler@typesafe.com)  
[polyglotprogramming.com/talks](http://polyglotprogramming.com/talks)  
[@deanwampler](https://twitter.com/deanwampler)

Tuesday, September 30, 14

About me. You can find this presentation and others on Big Data and Scala at [polyglotprogramming.com](http://polyglotprogramming.com).

Photo: Time lapse at night in Colorado.