



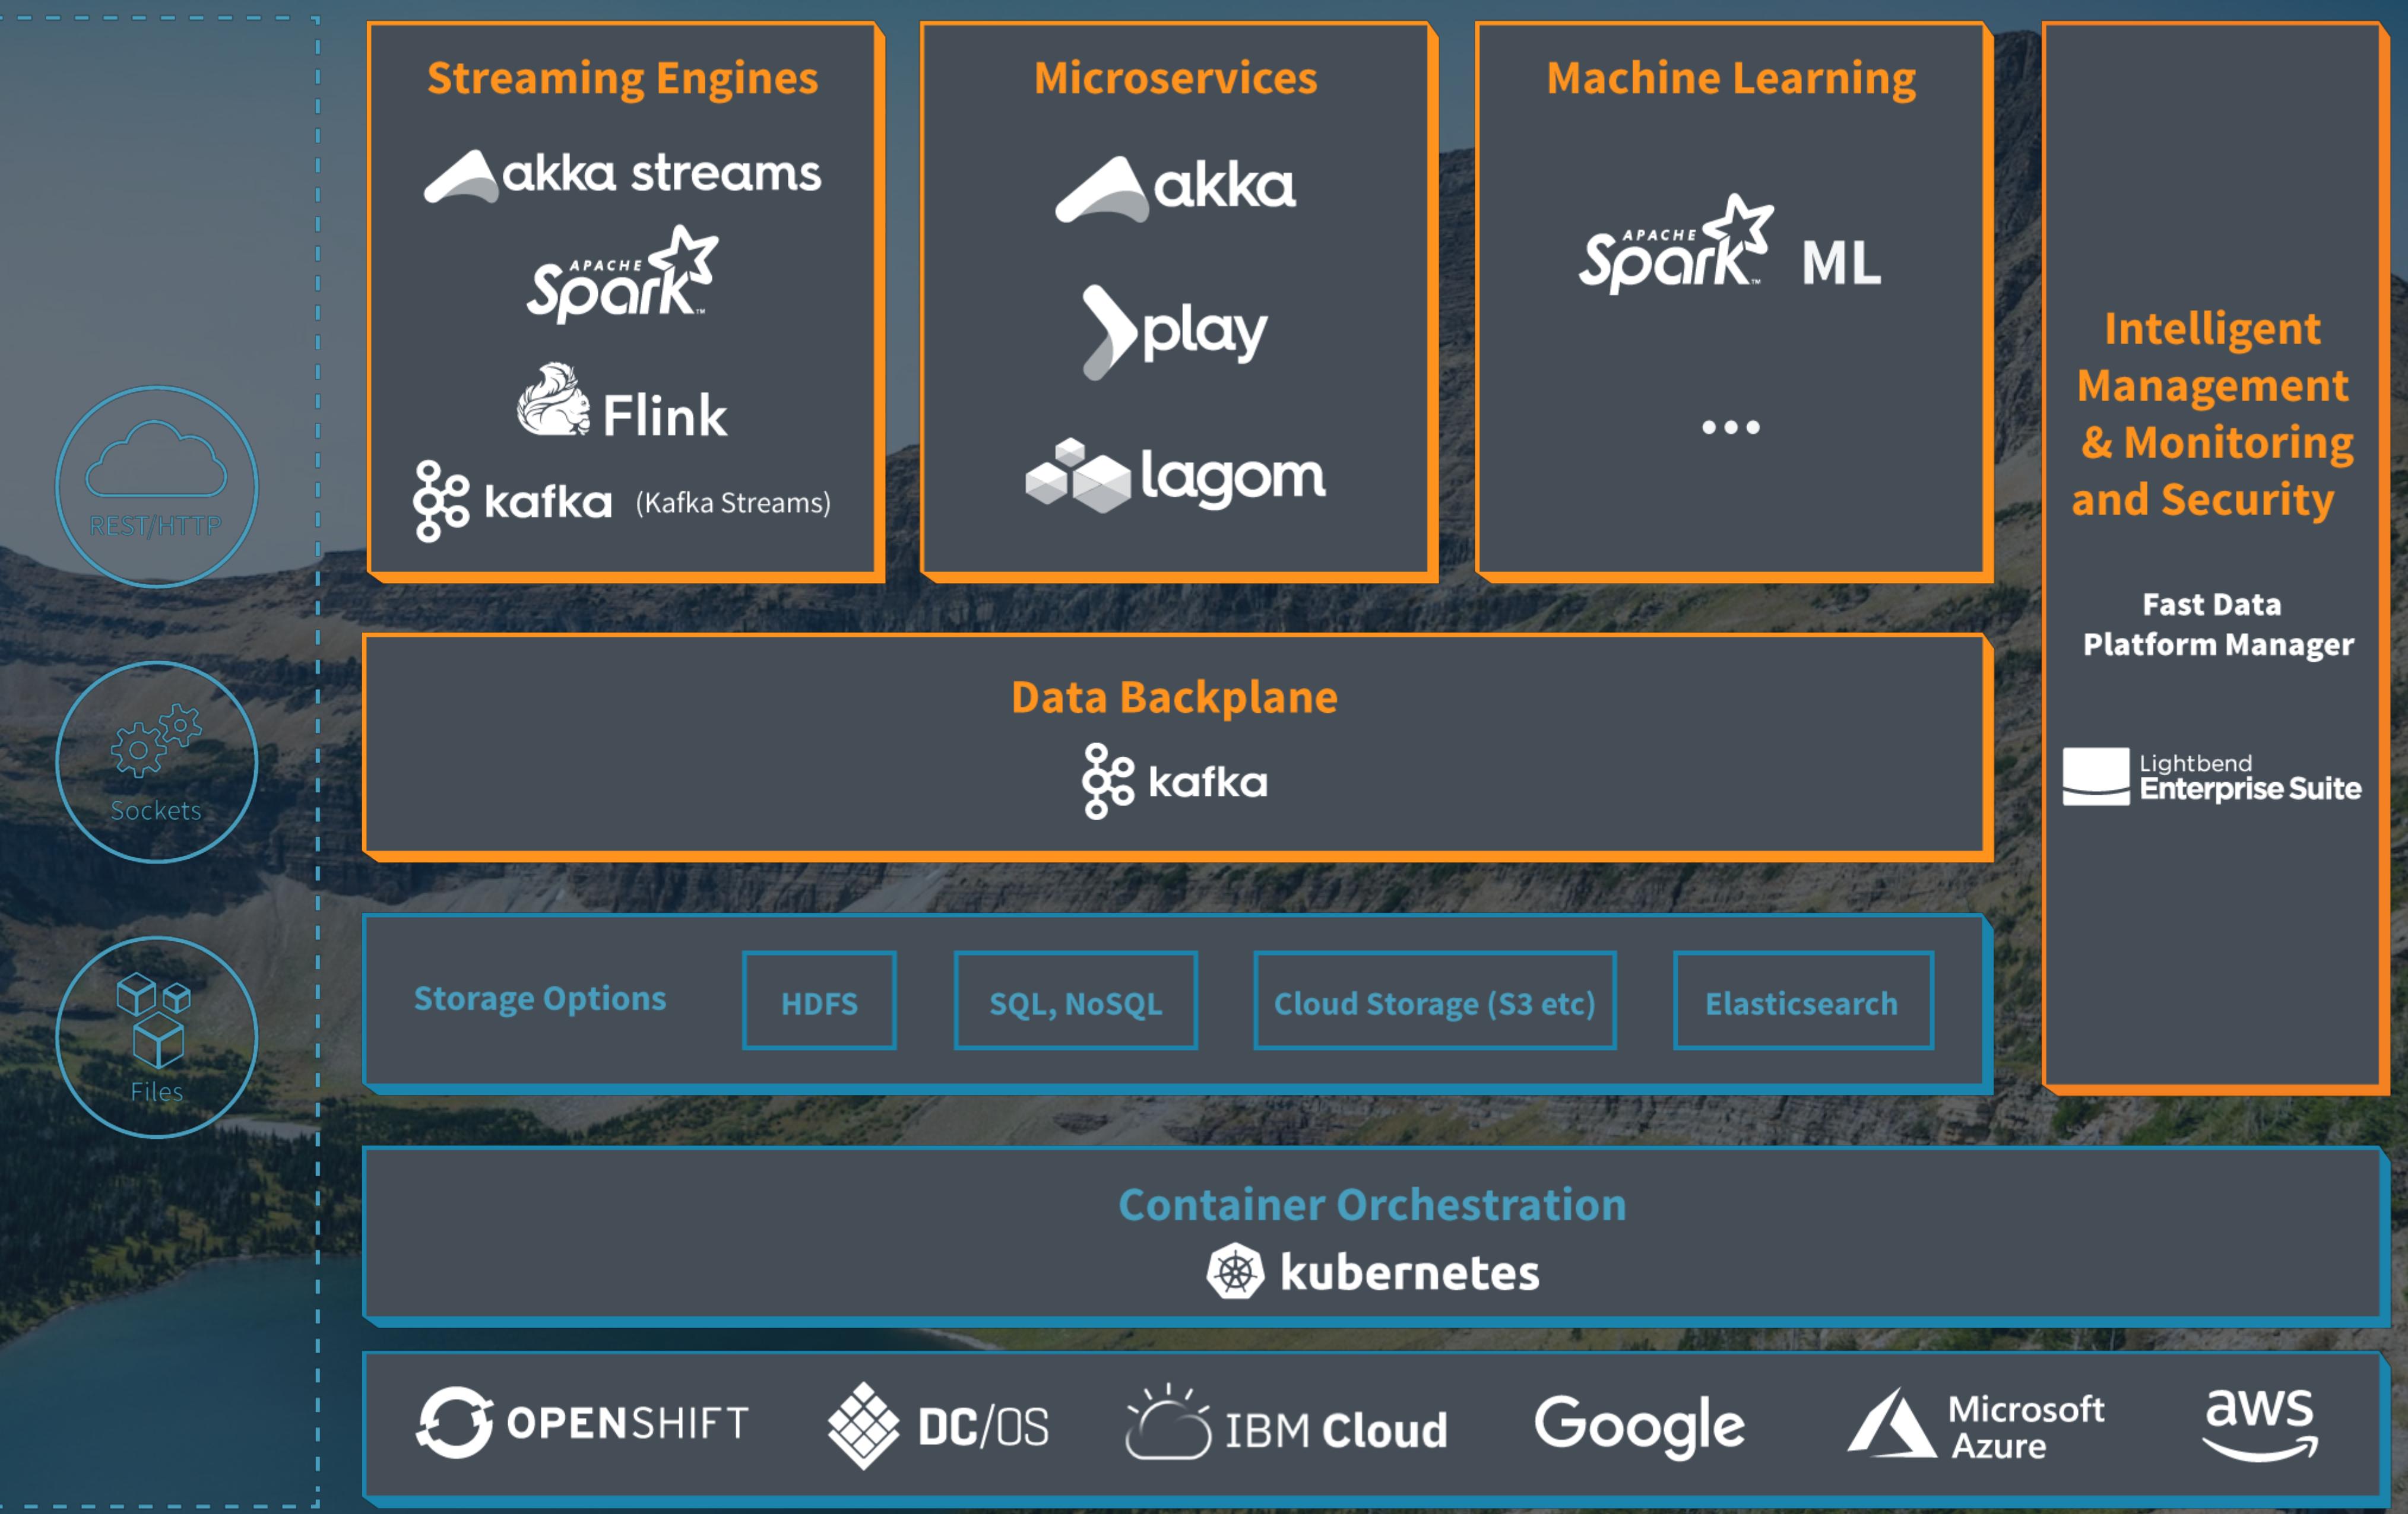
Streaming Data Microservices

With Akka Streams and Kafka Streams

Dean Wampler, Ph.D.
dean@lightbend.com
[@deanwampler](https://twitter.com/deanwampler)

Who am I?





I lead the Lightbend Fast Data Platform; unified streaming data & microservices

lightbend.com/fast-data-platform

Streaming Engines

akka streams



kafka (Kafka Streams)



Microservices



...

Machine Learning



Intelligent
Management
& Monitoring
and Security

Fast Data
Platform Manager

Lightbend
Enterprise Suite

Data Backplane



Storage Options

HDFS

SQL, NoSQL

Cloud Storage (S3 etc)

Elasticsearch



Container Orchestration





Free as in 

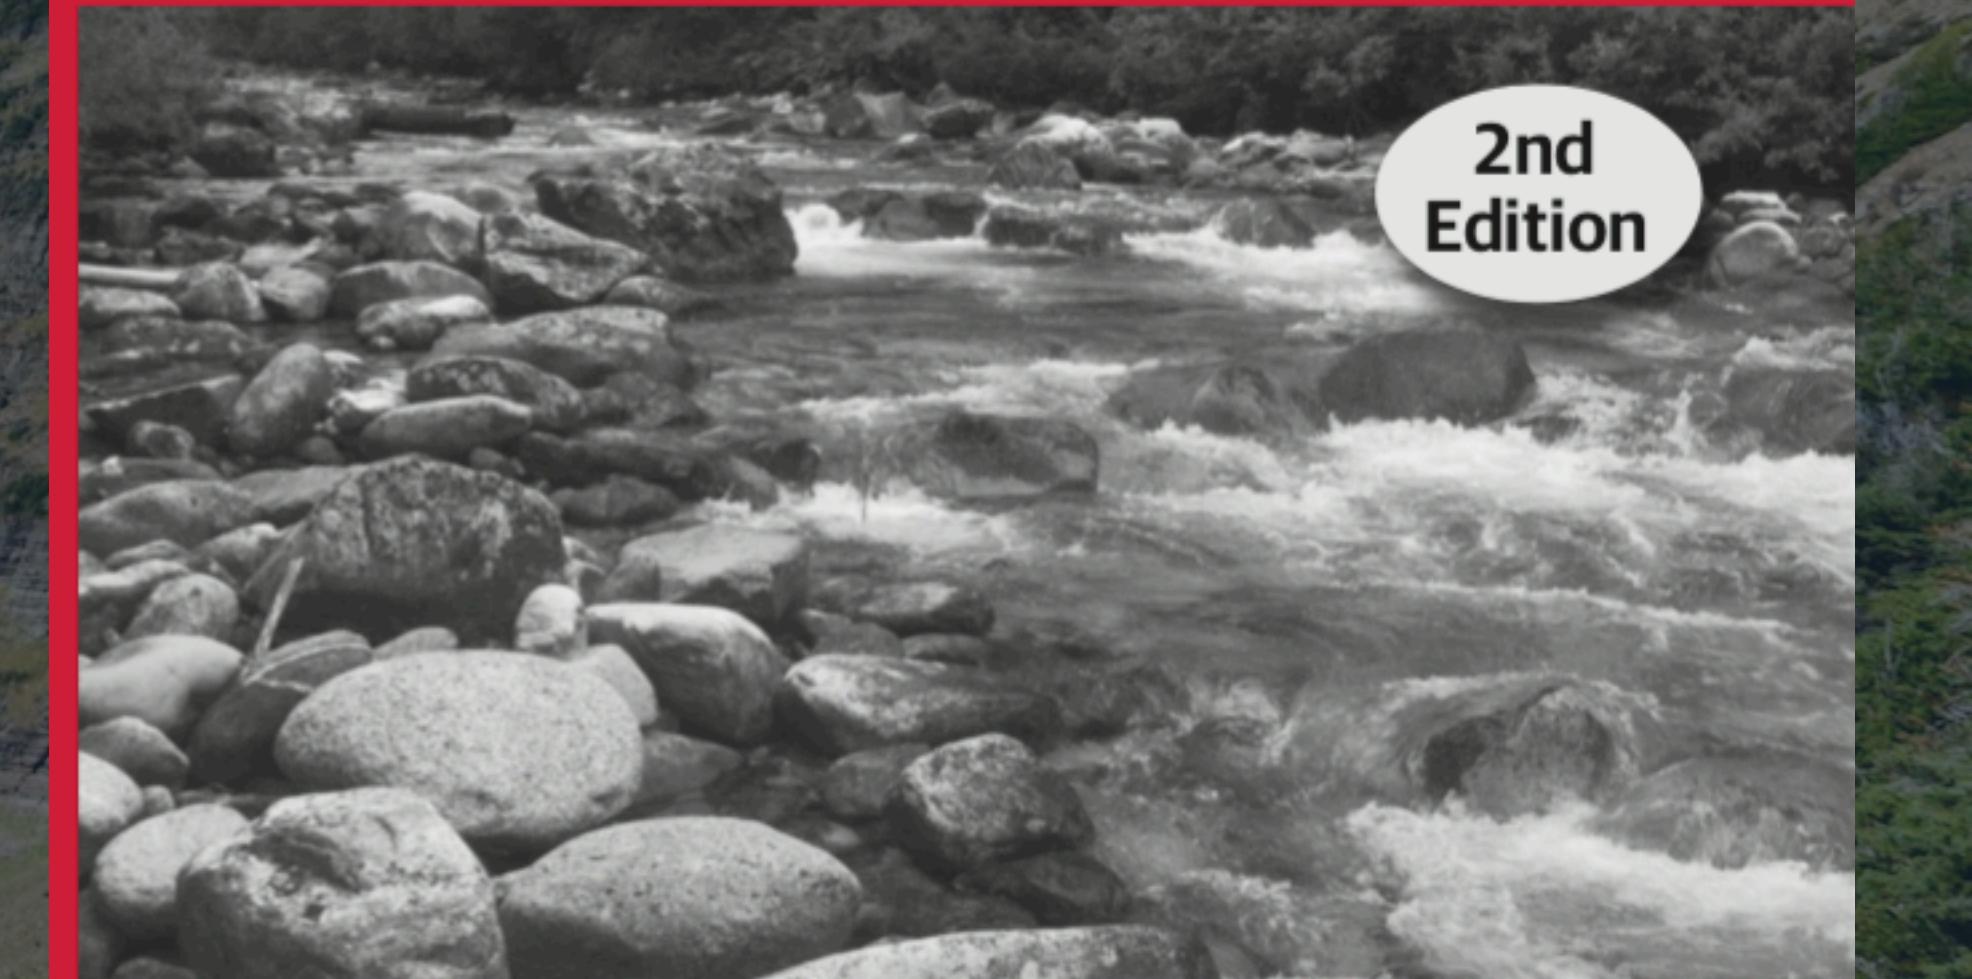
New second edition!
lbnd.io/fast-data-book

O'REILLY®

Compliments of
 **Lightbend**

Fast Data Architectures for Streaming Applications

Getting Answers Now from
Data Sets That Never End

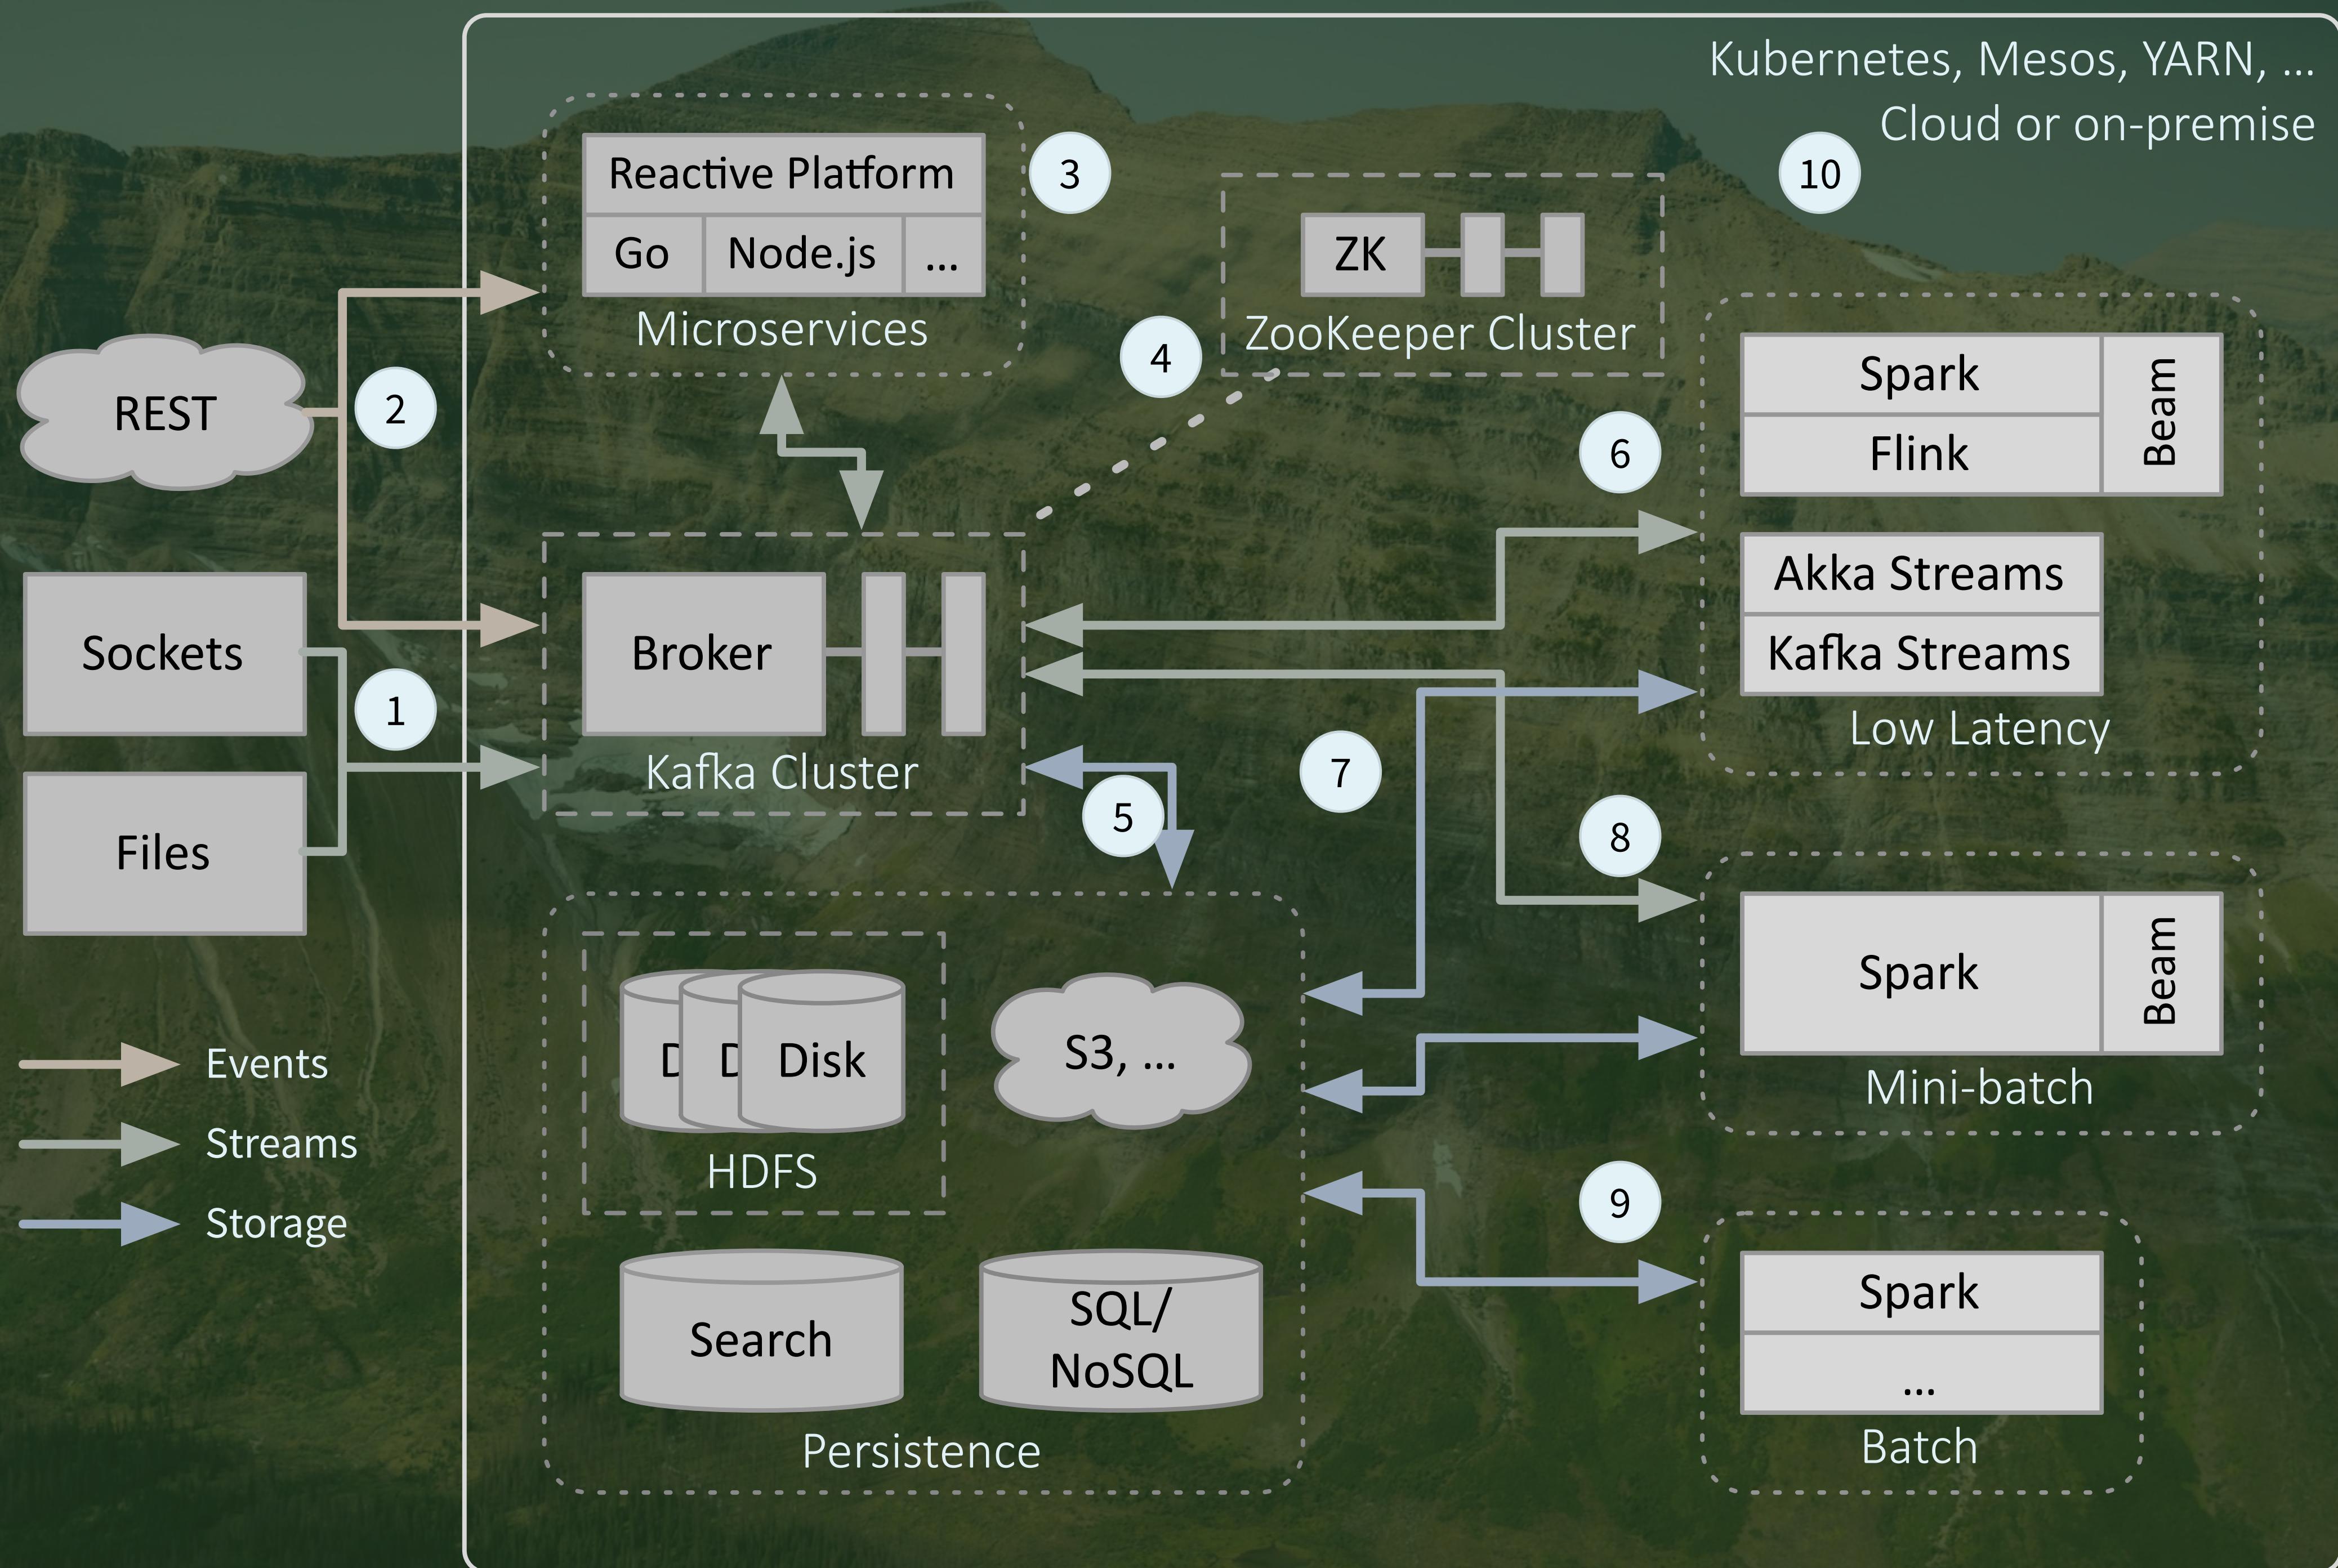


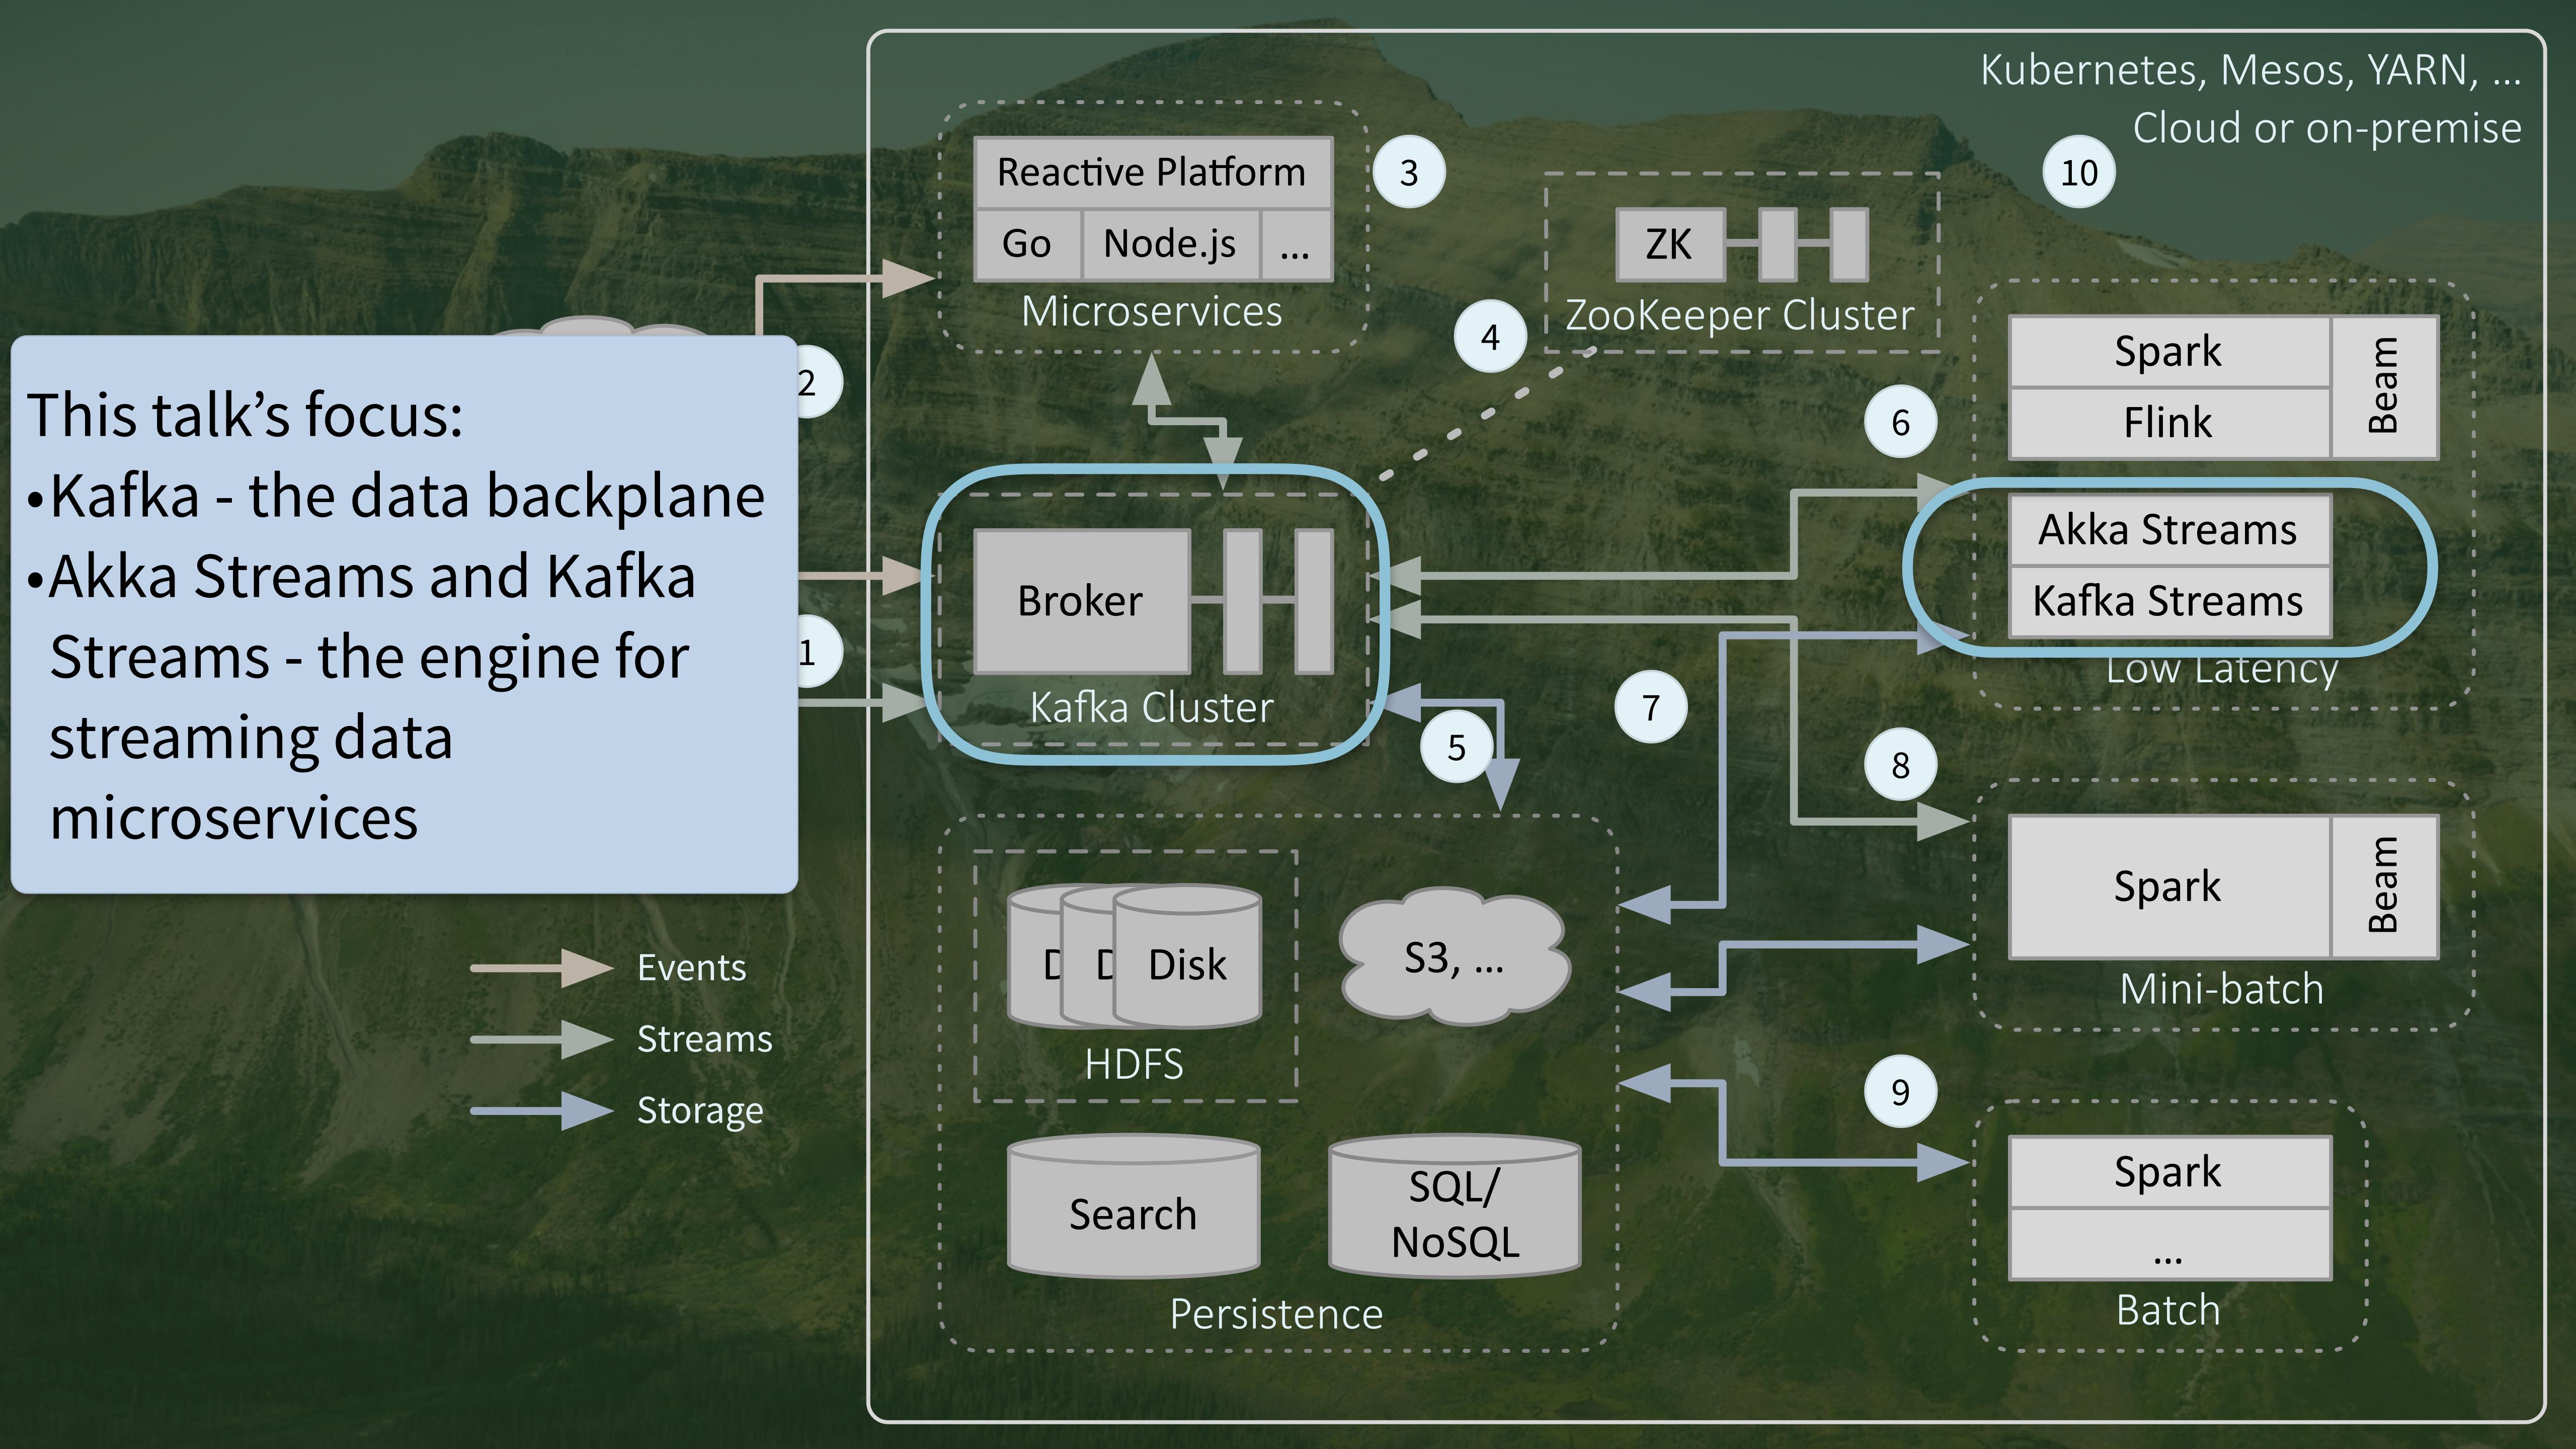
2nd
Edition

Dean Wampler, PhD



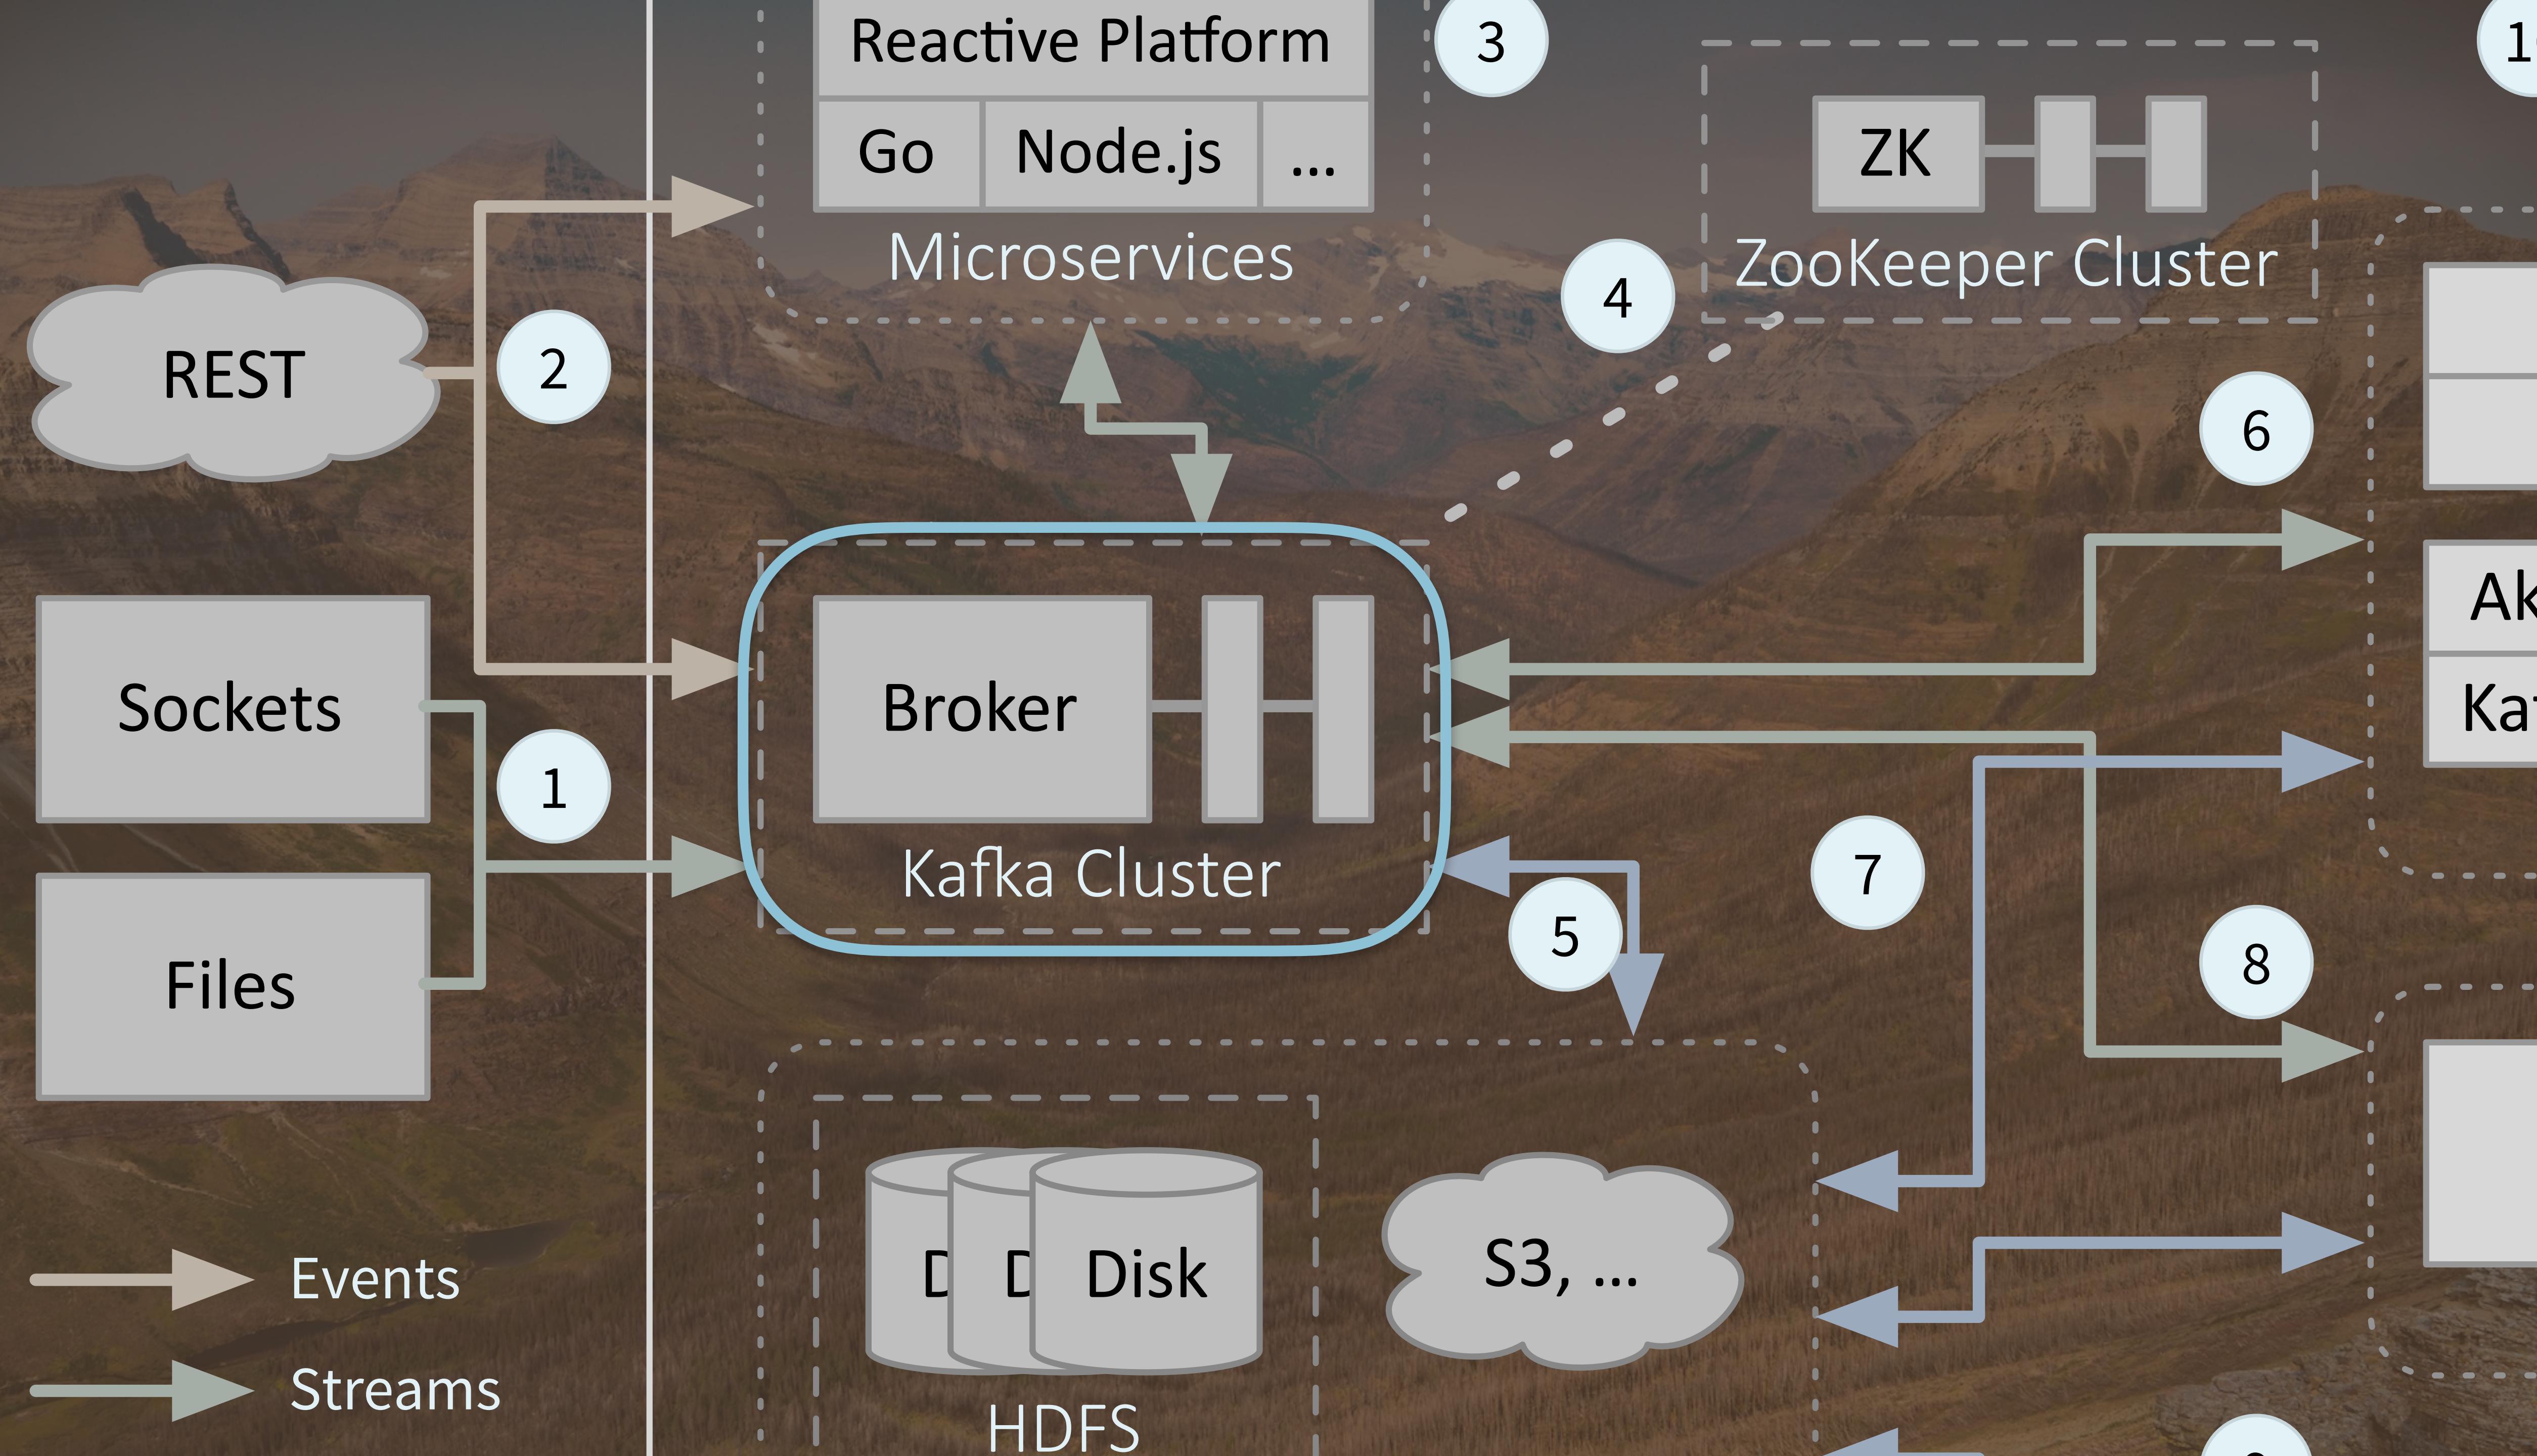
Streaming architectures (from the report)

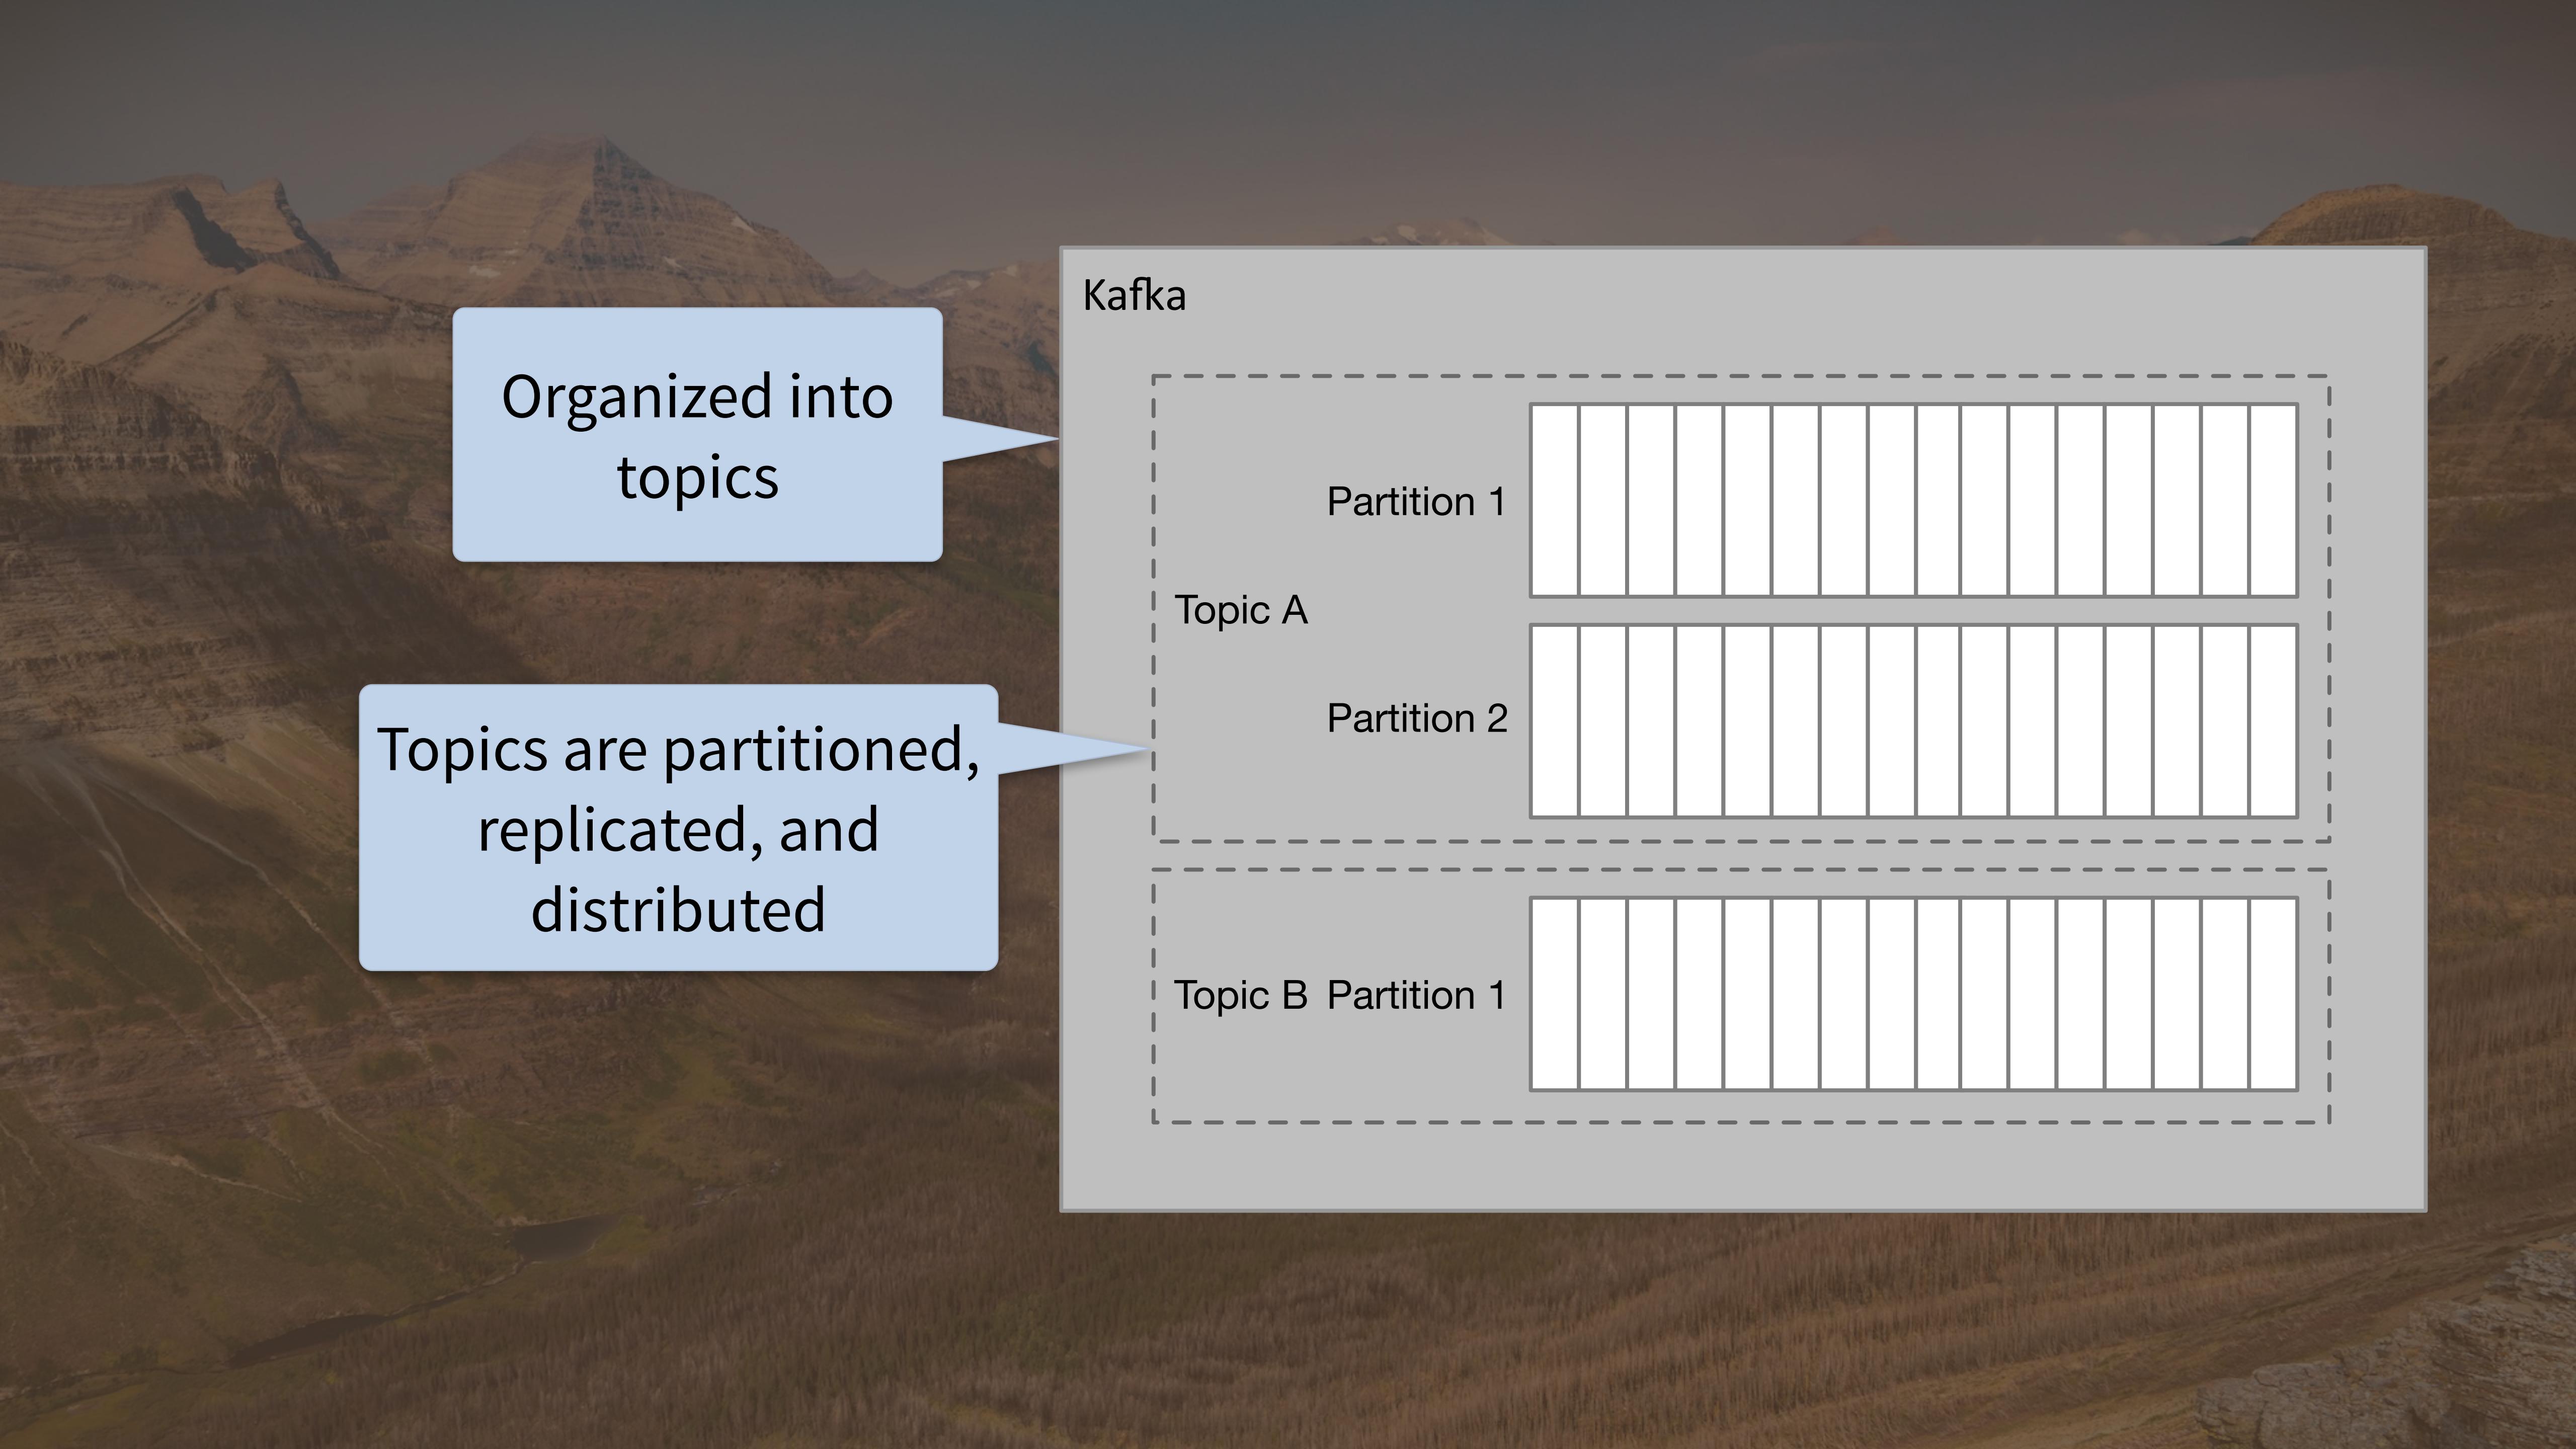




A wide-angle photograph of the Rocky Mountains. In the foreground, a deep valley is filled with a dense forest of green trees, with a dark river or stream winding its way through it. The middle ground shows the steep, rocky slopes of the mountains, which are partially covered in sparse vegetation and patches of snow. The background features a range of mountains with prominent peaks, some of which are capped with white snow. The sky is clear and blue.

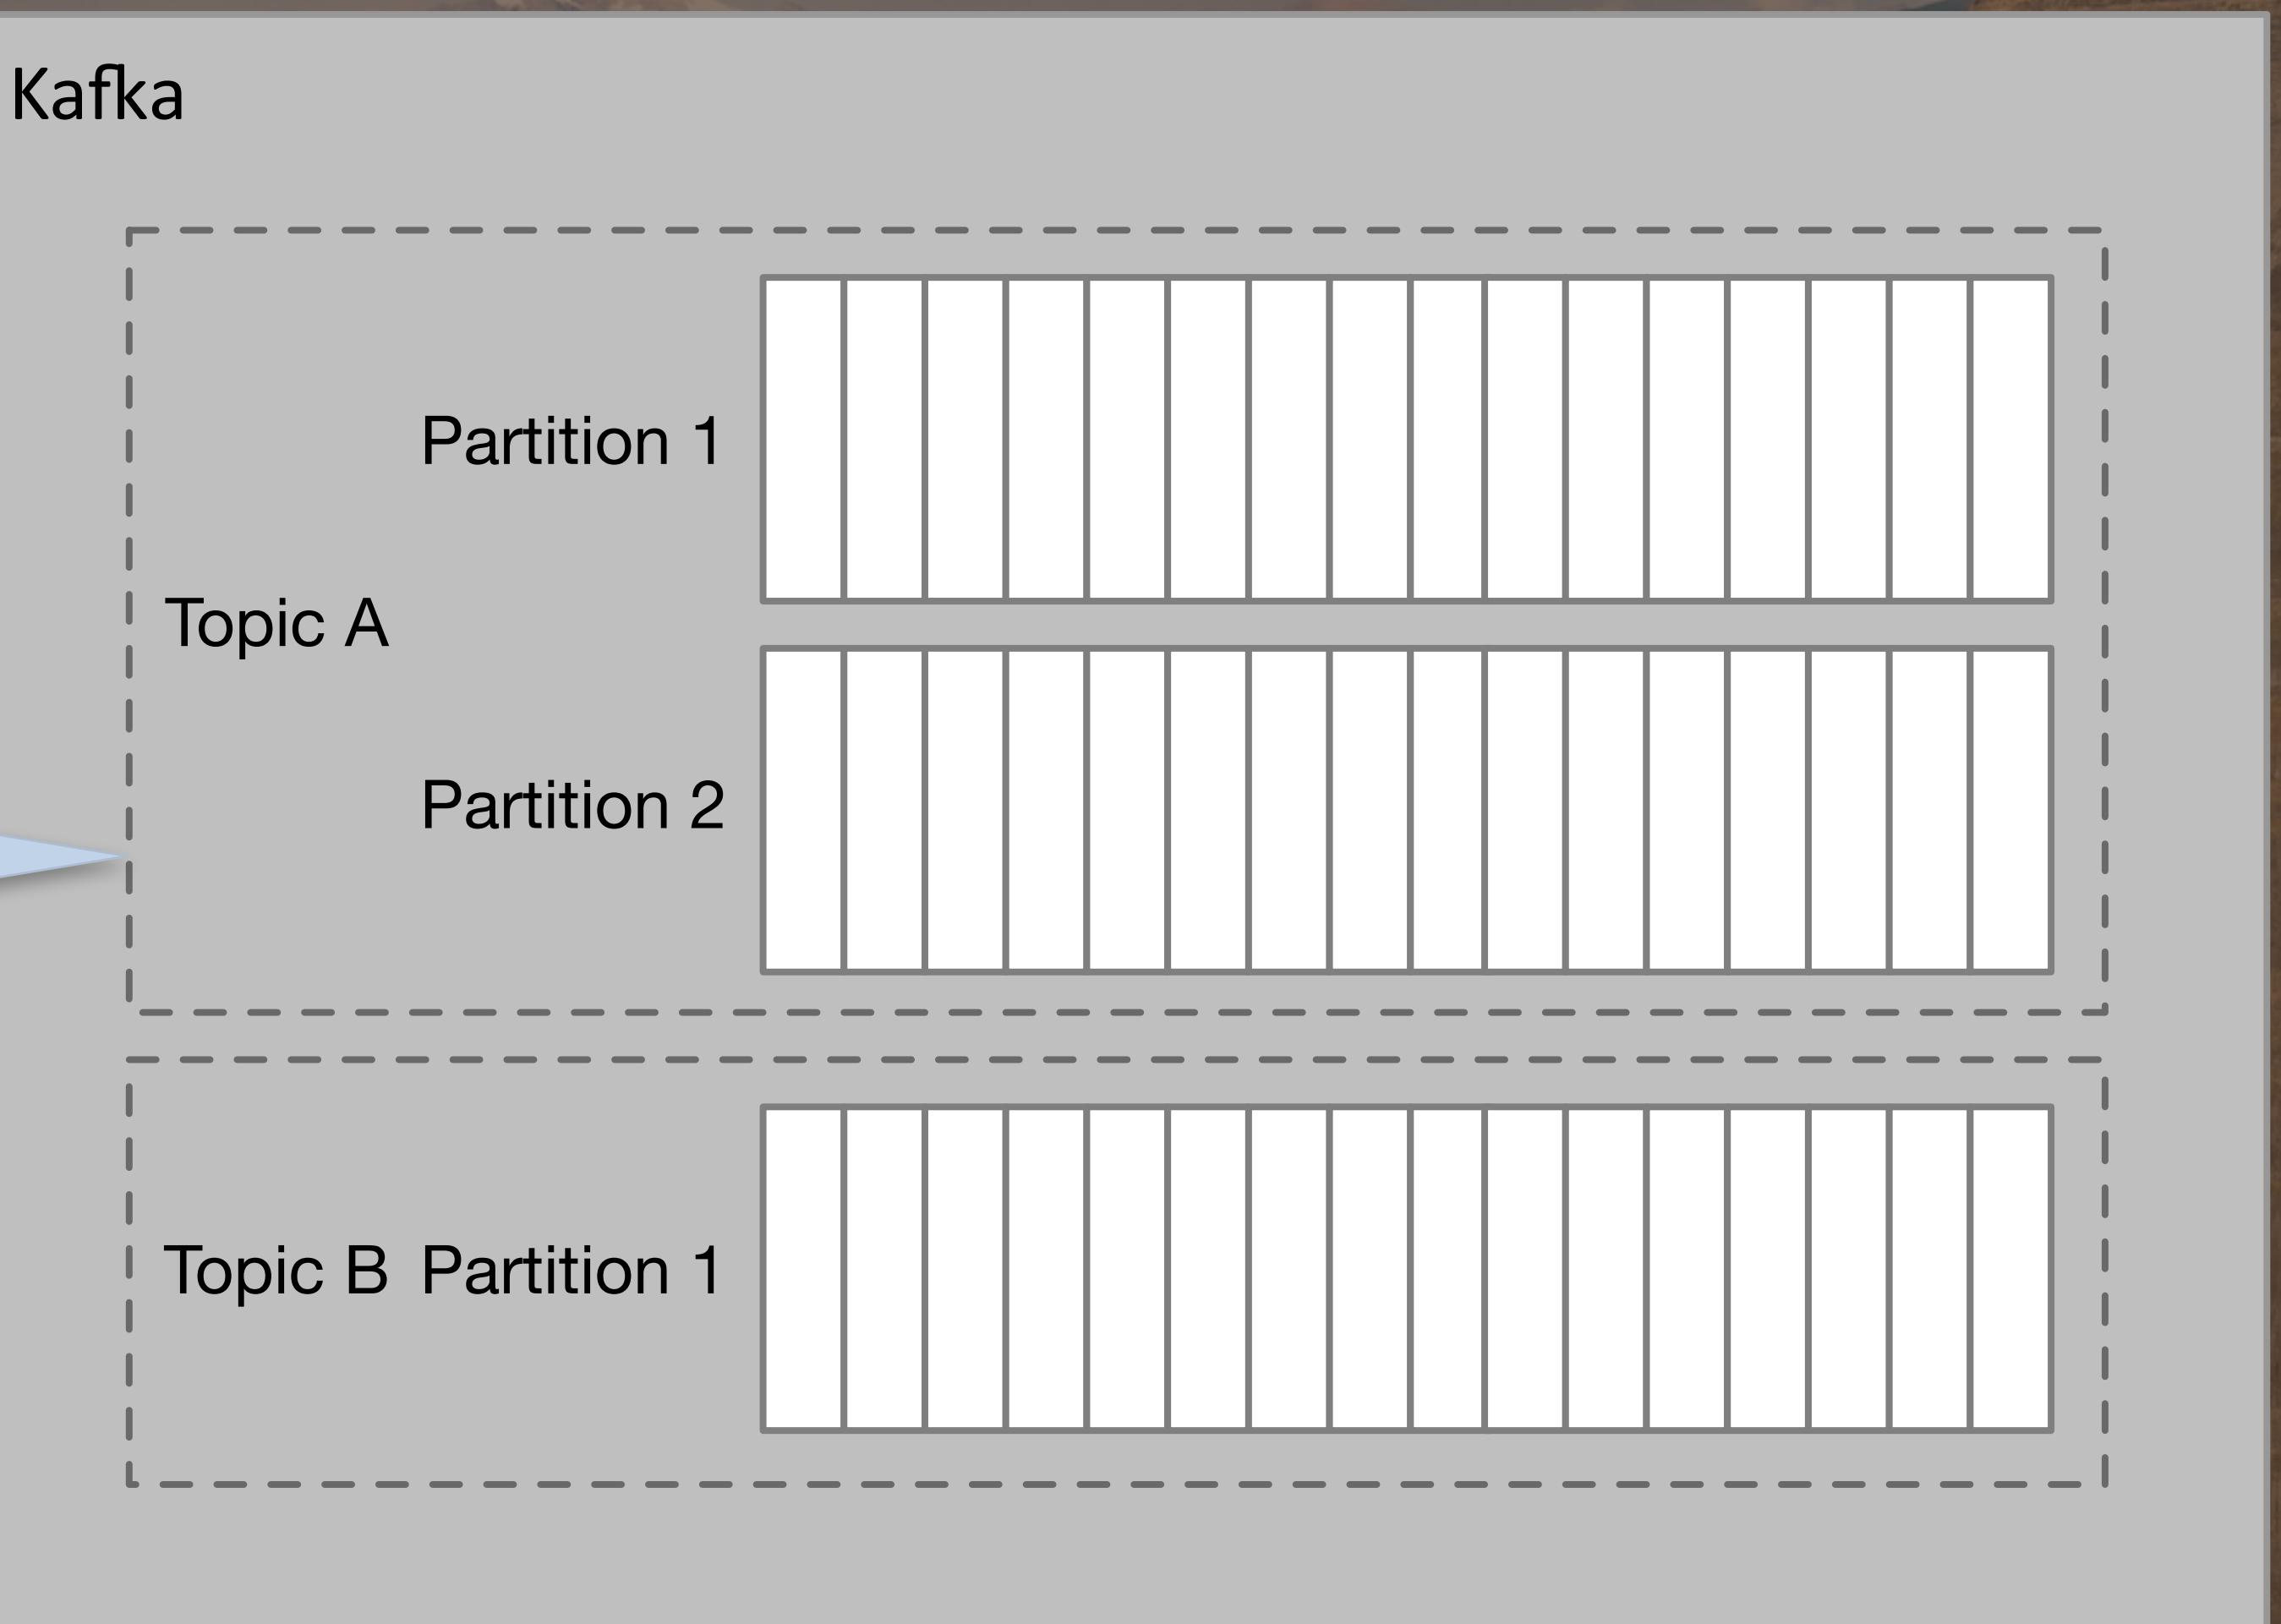
Why Kafka?



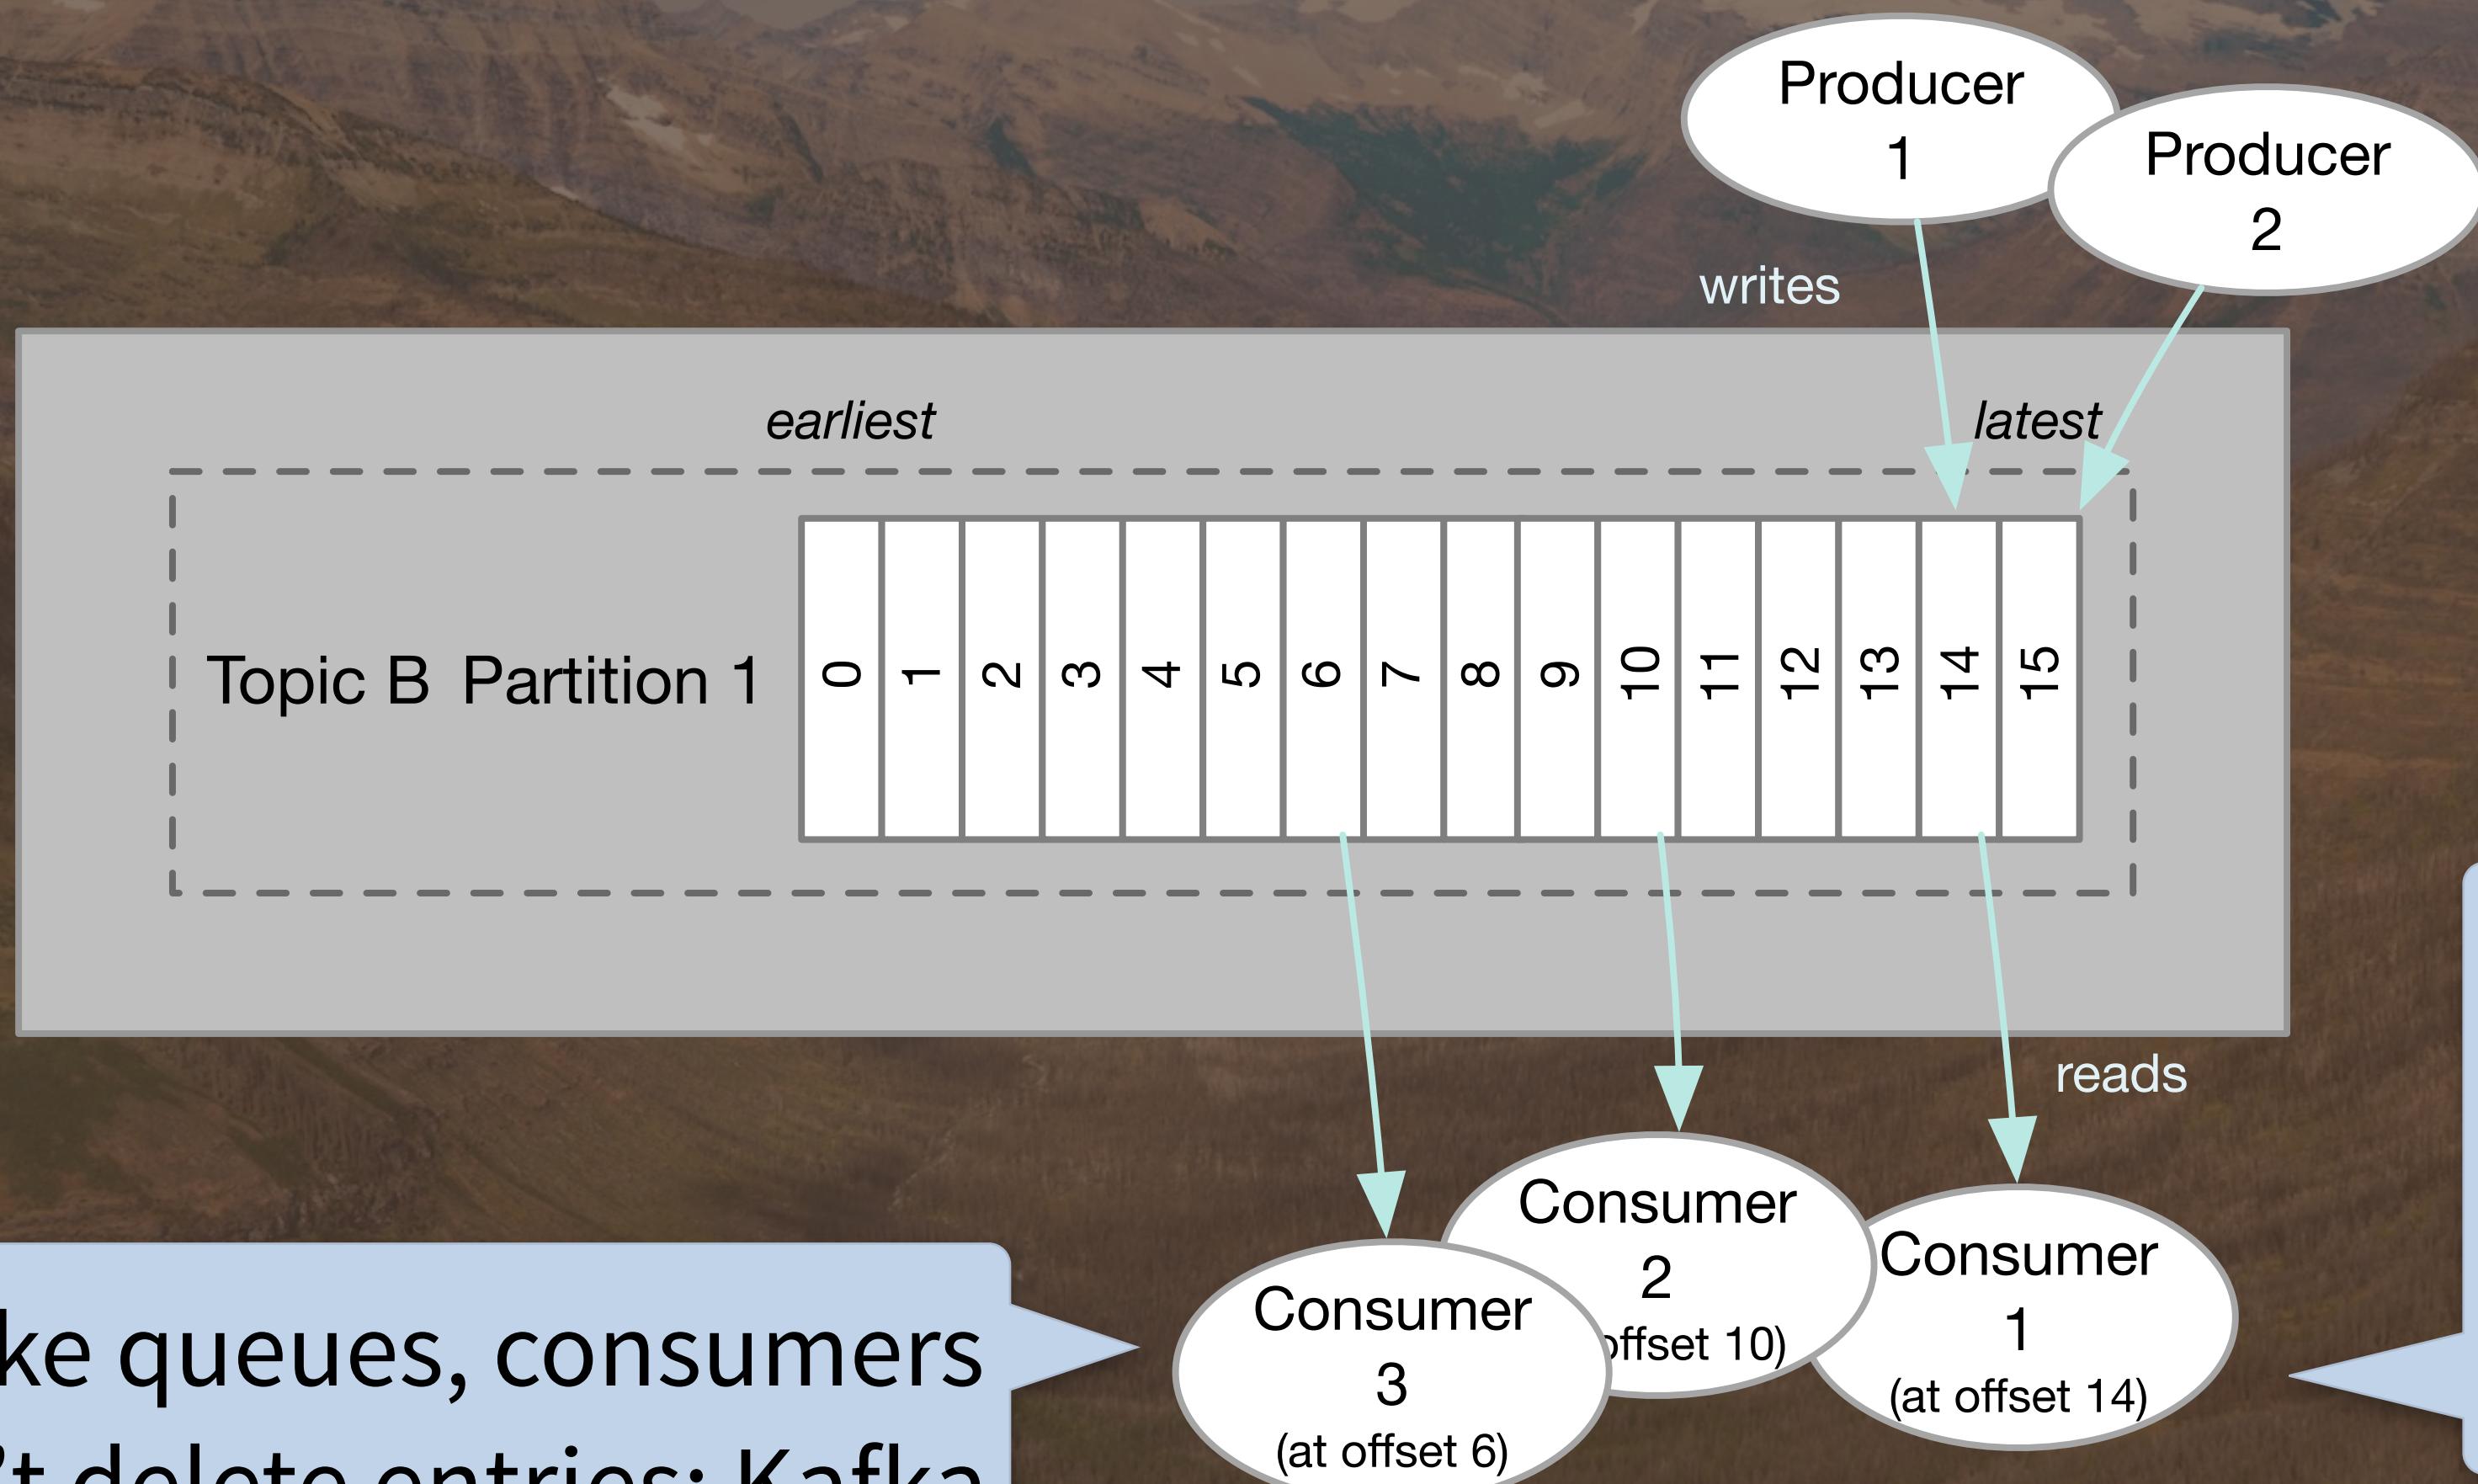


Organized into topics

Topics are partitioned,
replicated, and
distributed

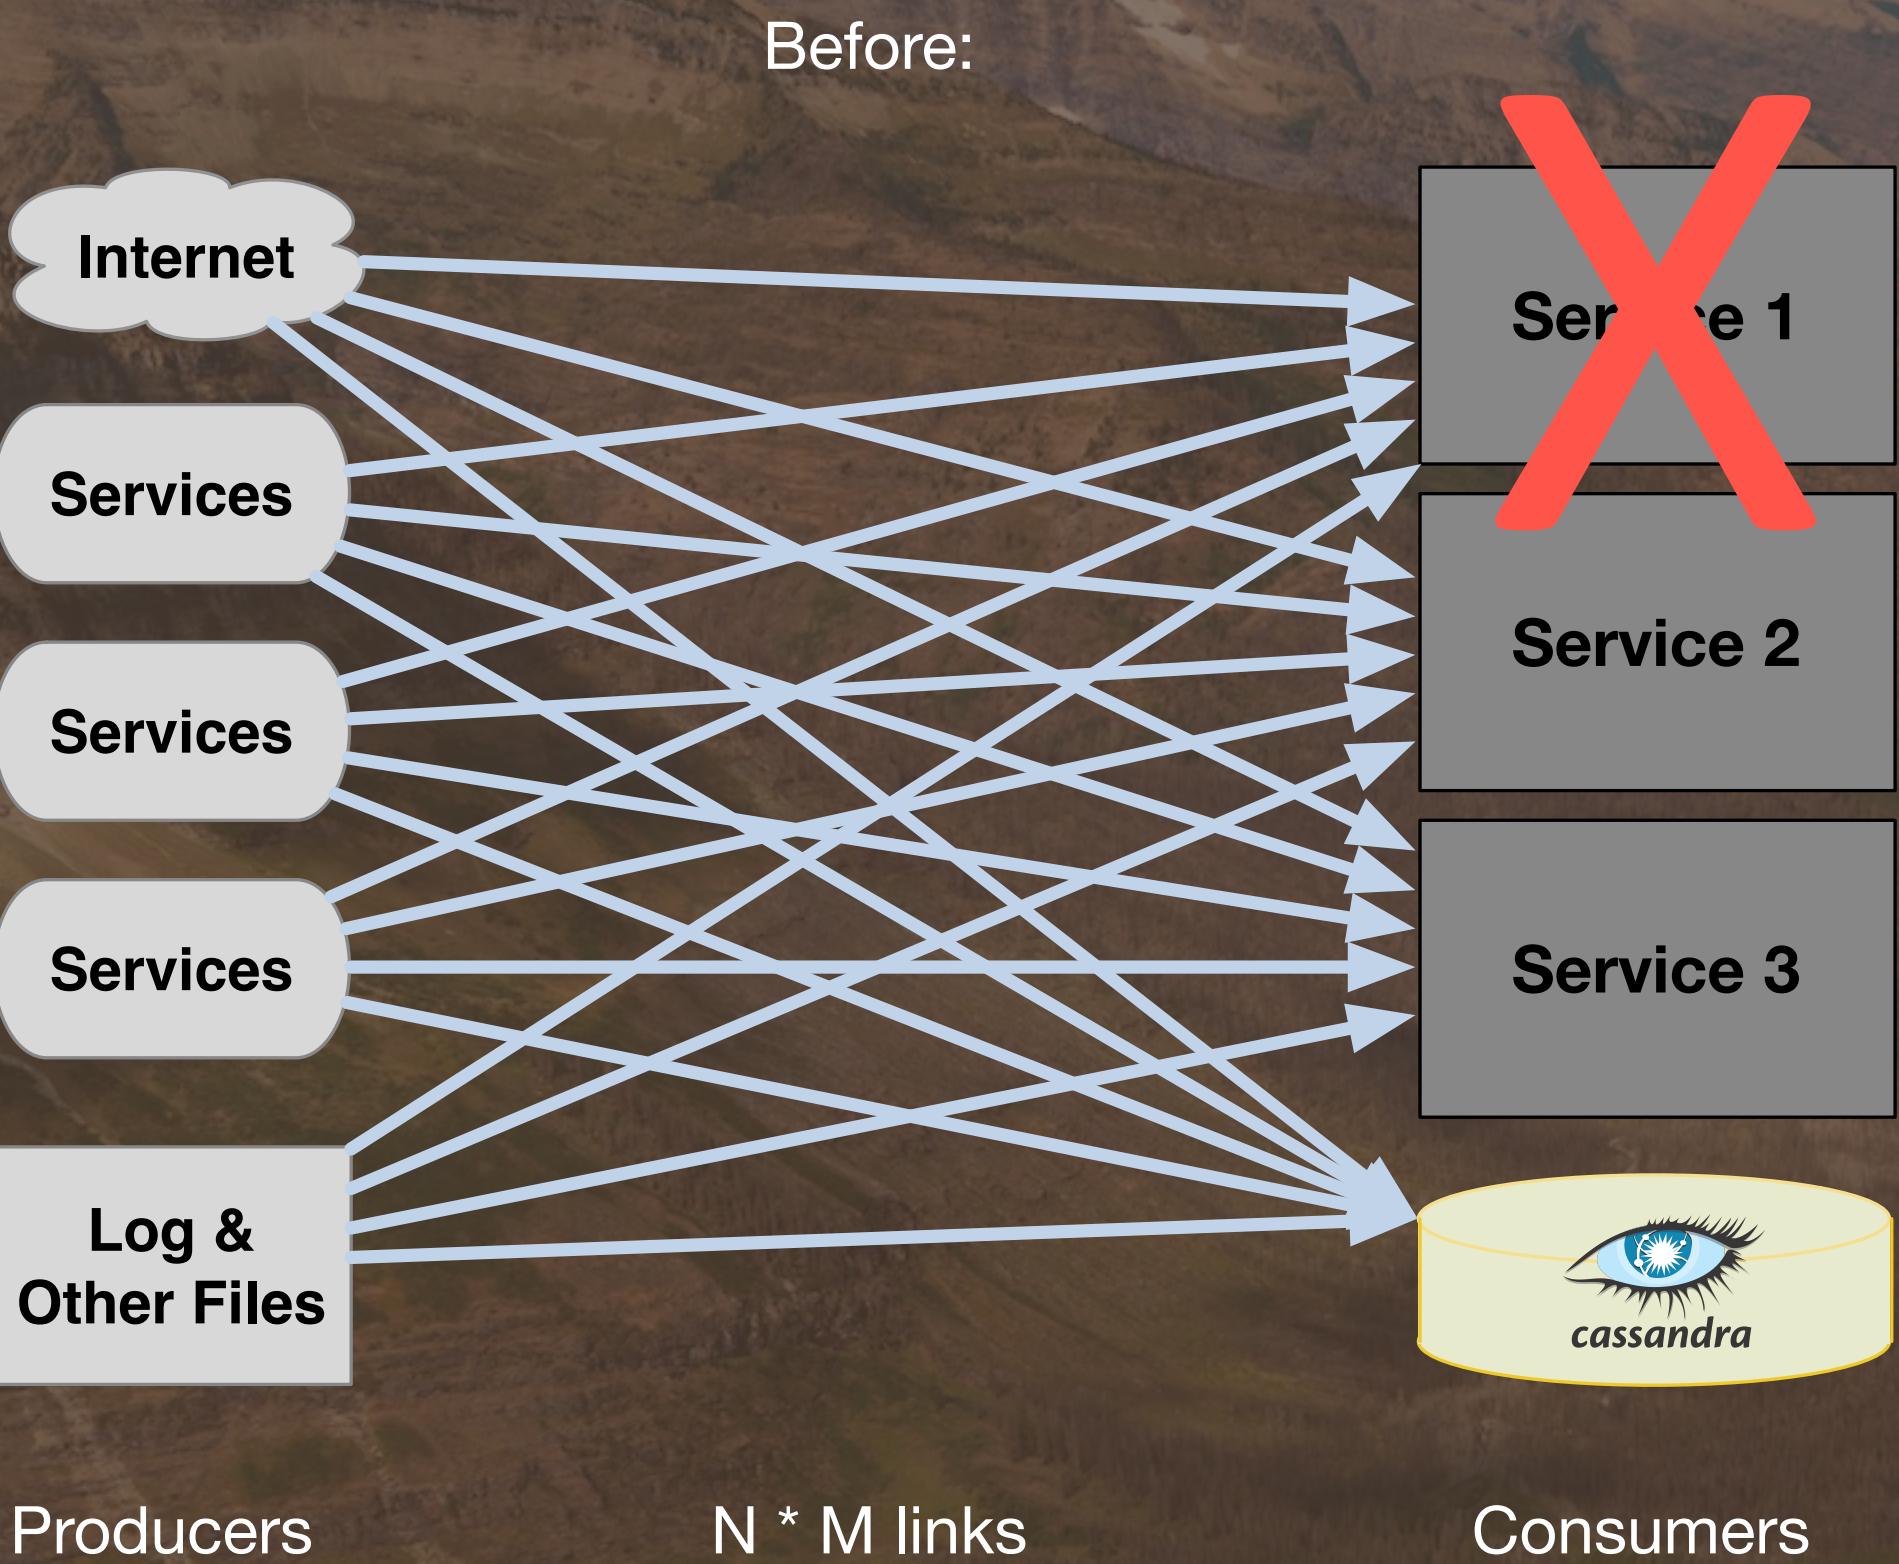


Logs, not queues!

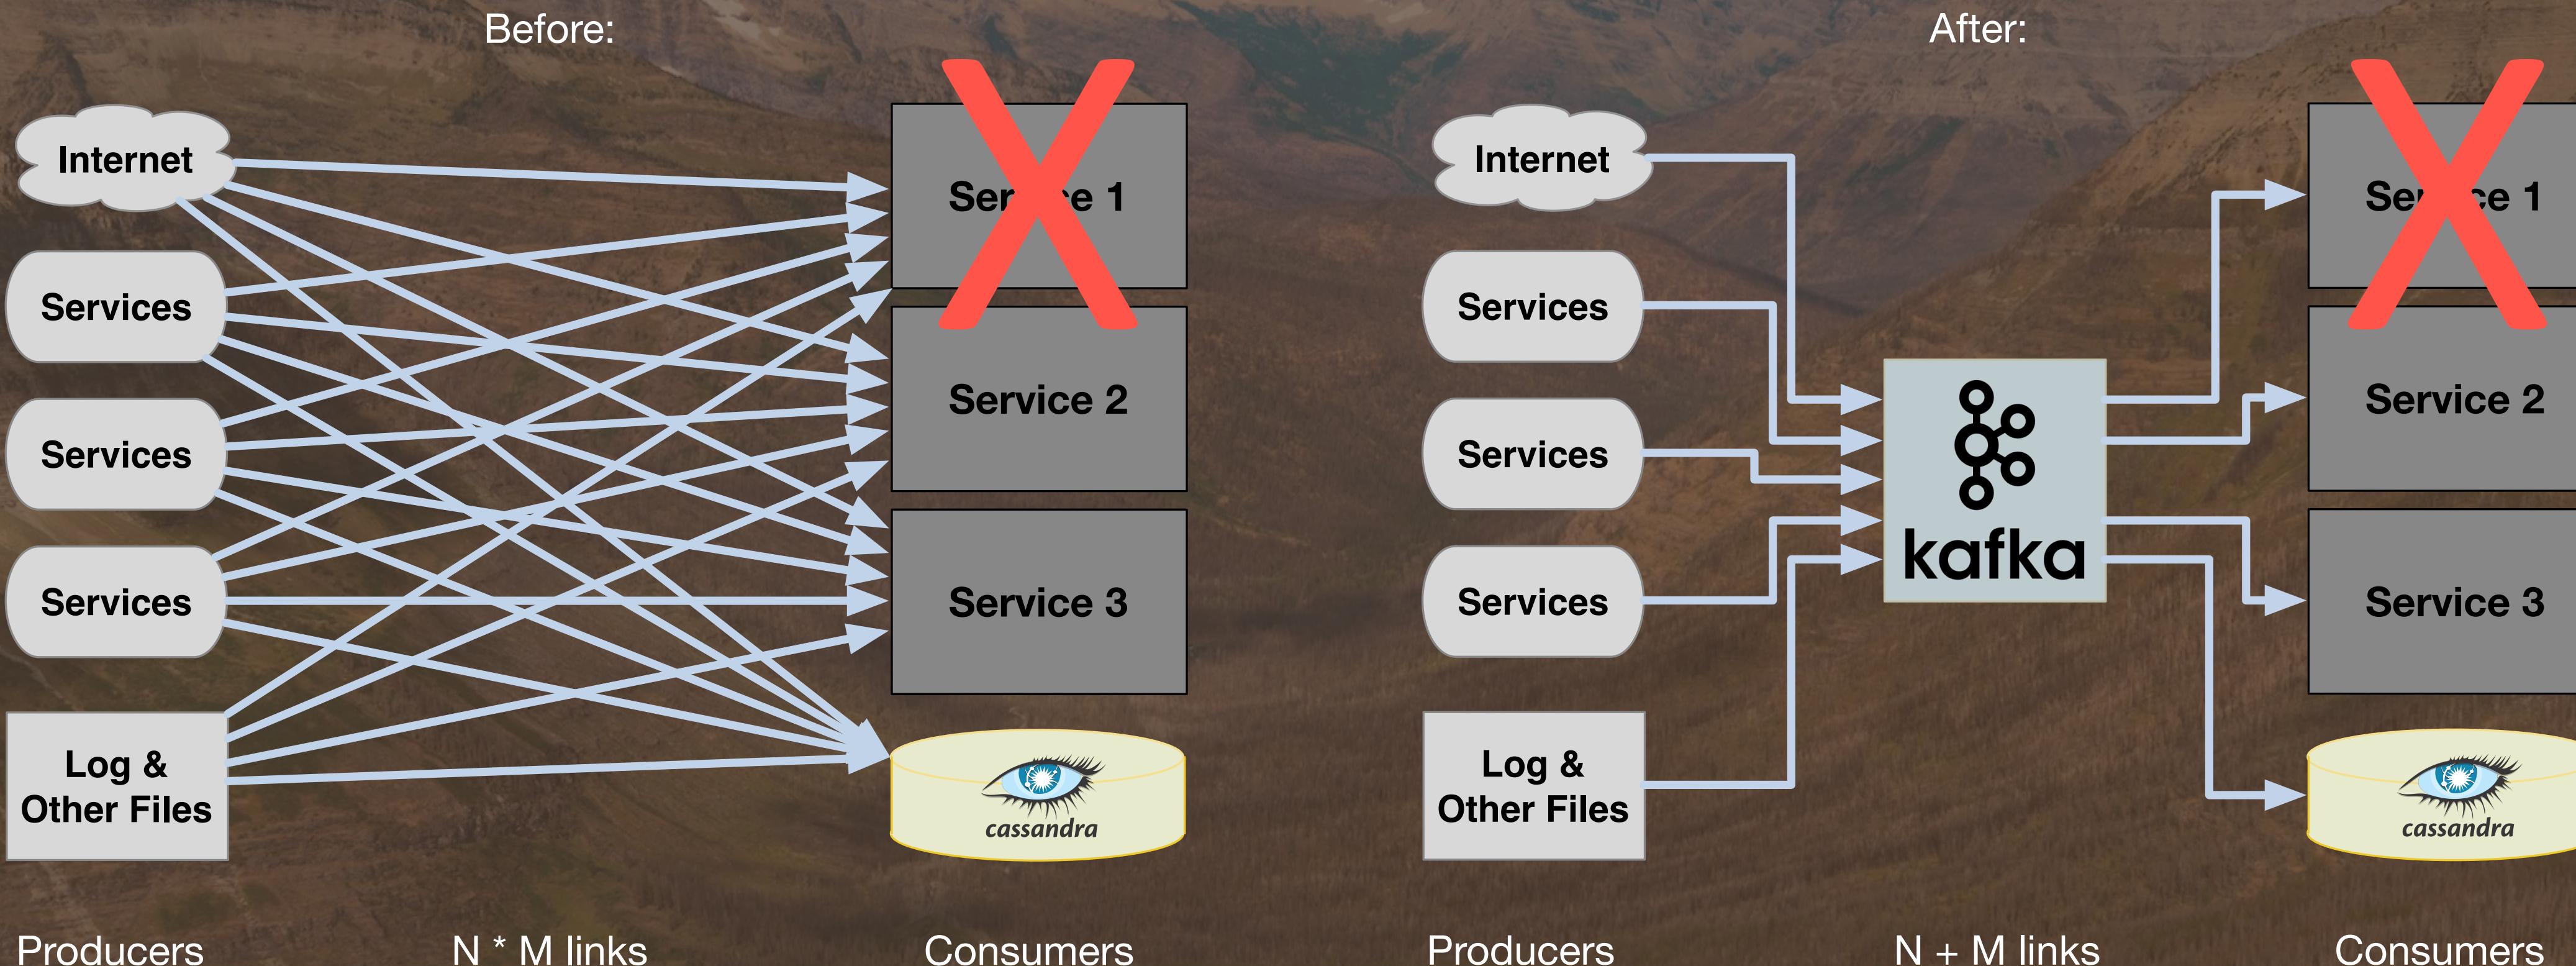


Unlike queues, consumers don't delete entries; Kafka manages their lifecycles

Architectural Benefits of Kafka

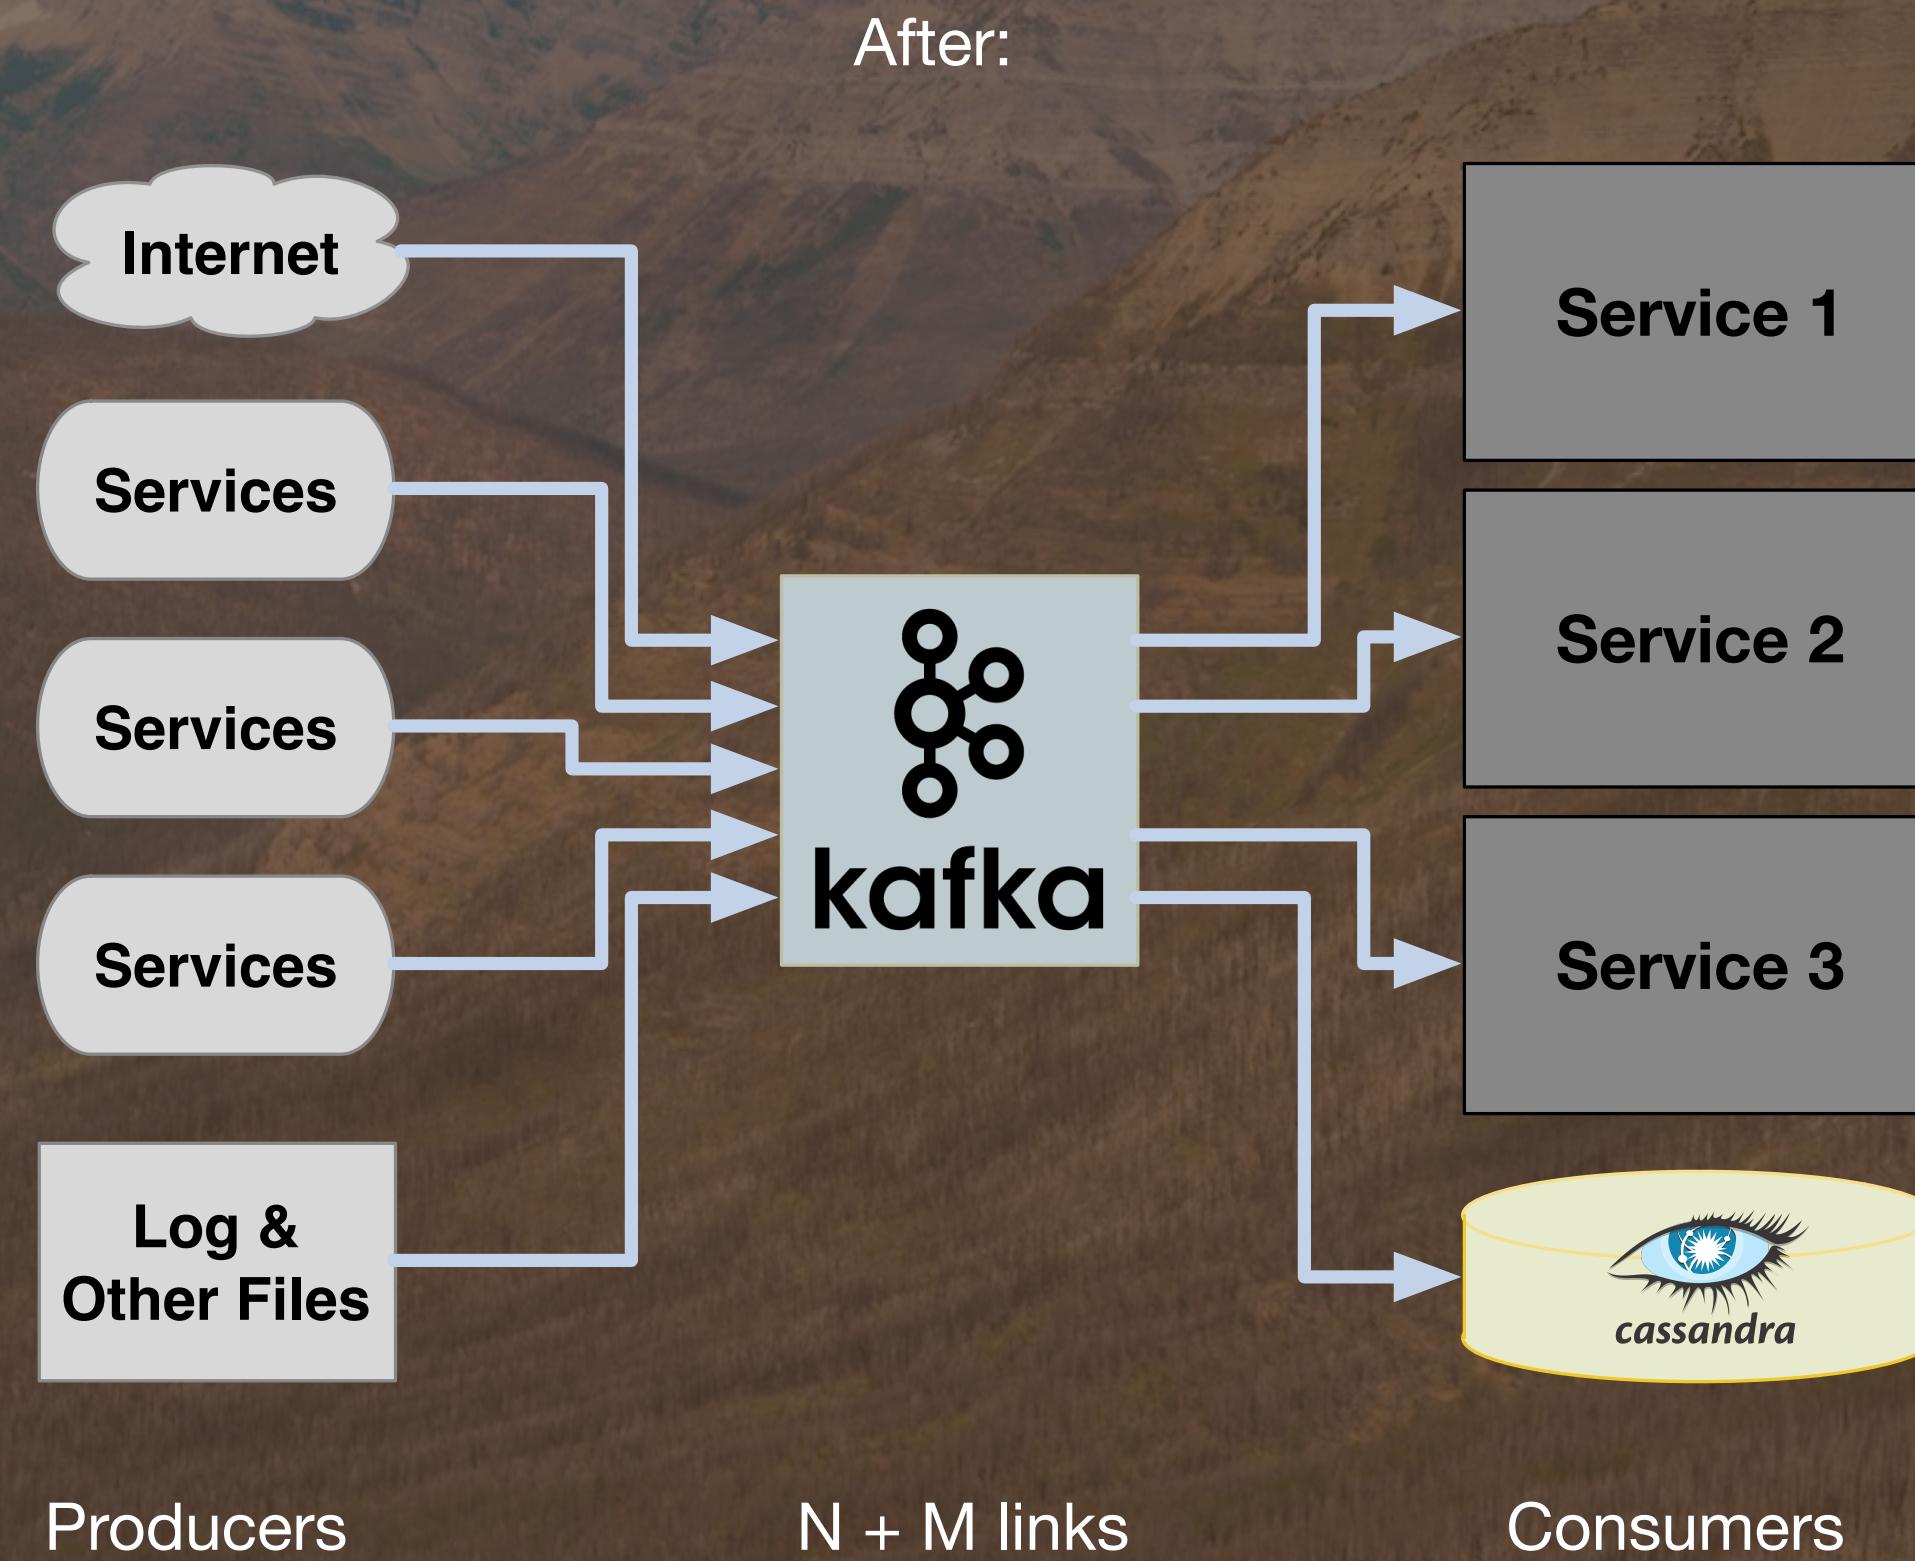


Architectural Benefits of Kafka



Architectural Benefits of Kafka

- Simplify dependencies
- Resilient against data loss
- M producers, N consumers
- Simplicity of one “API” for communication

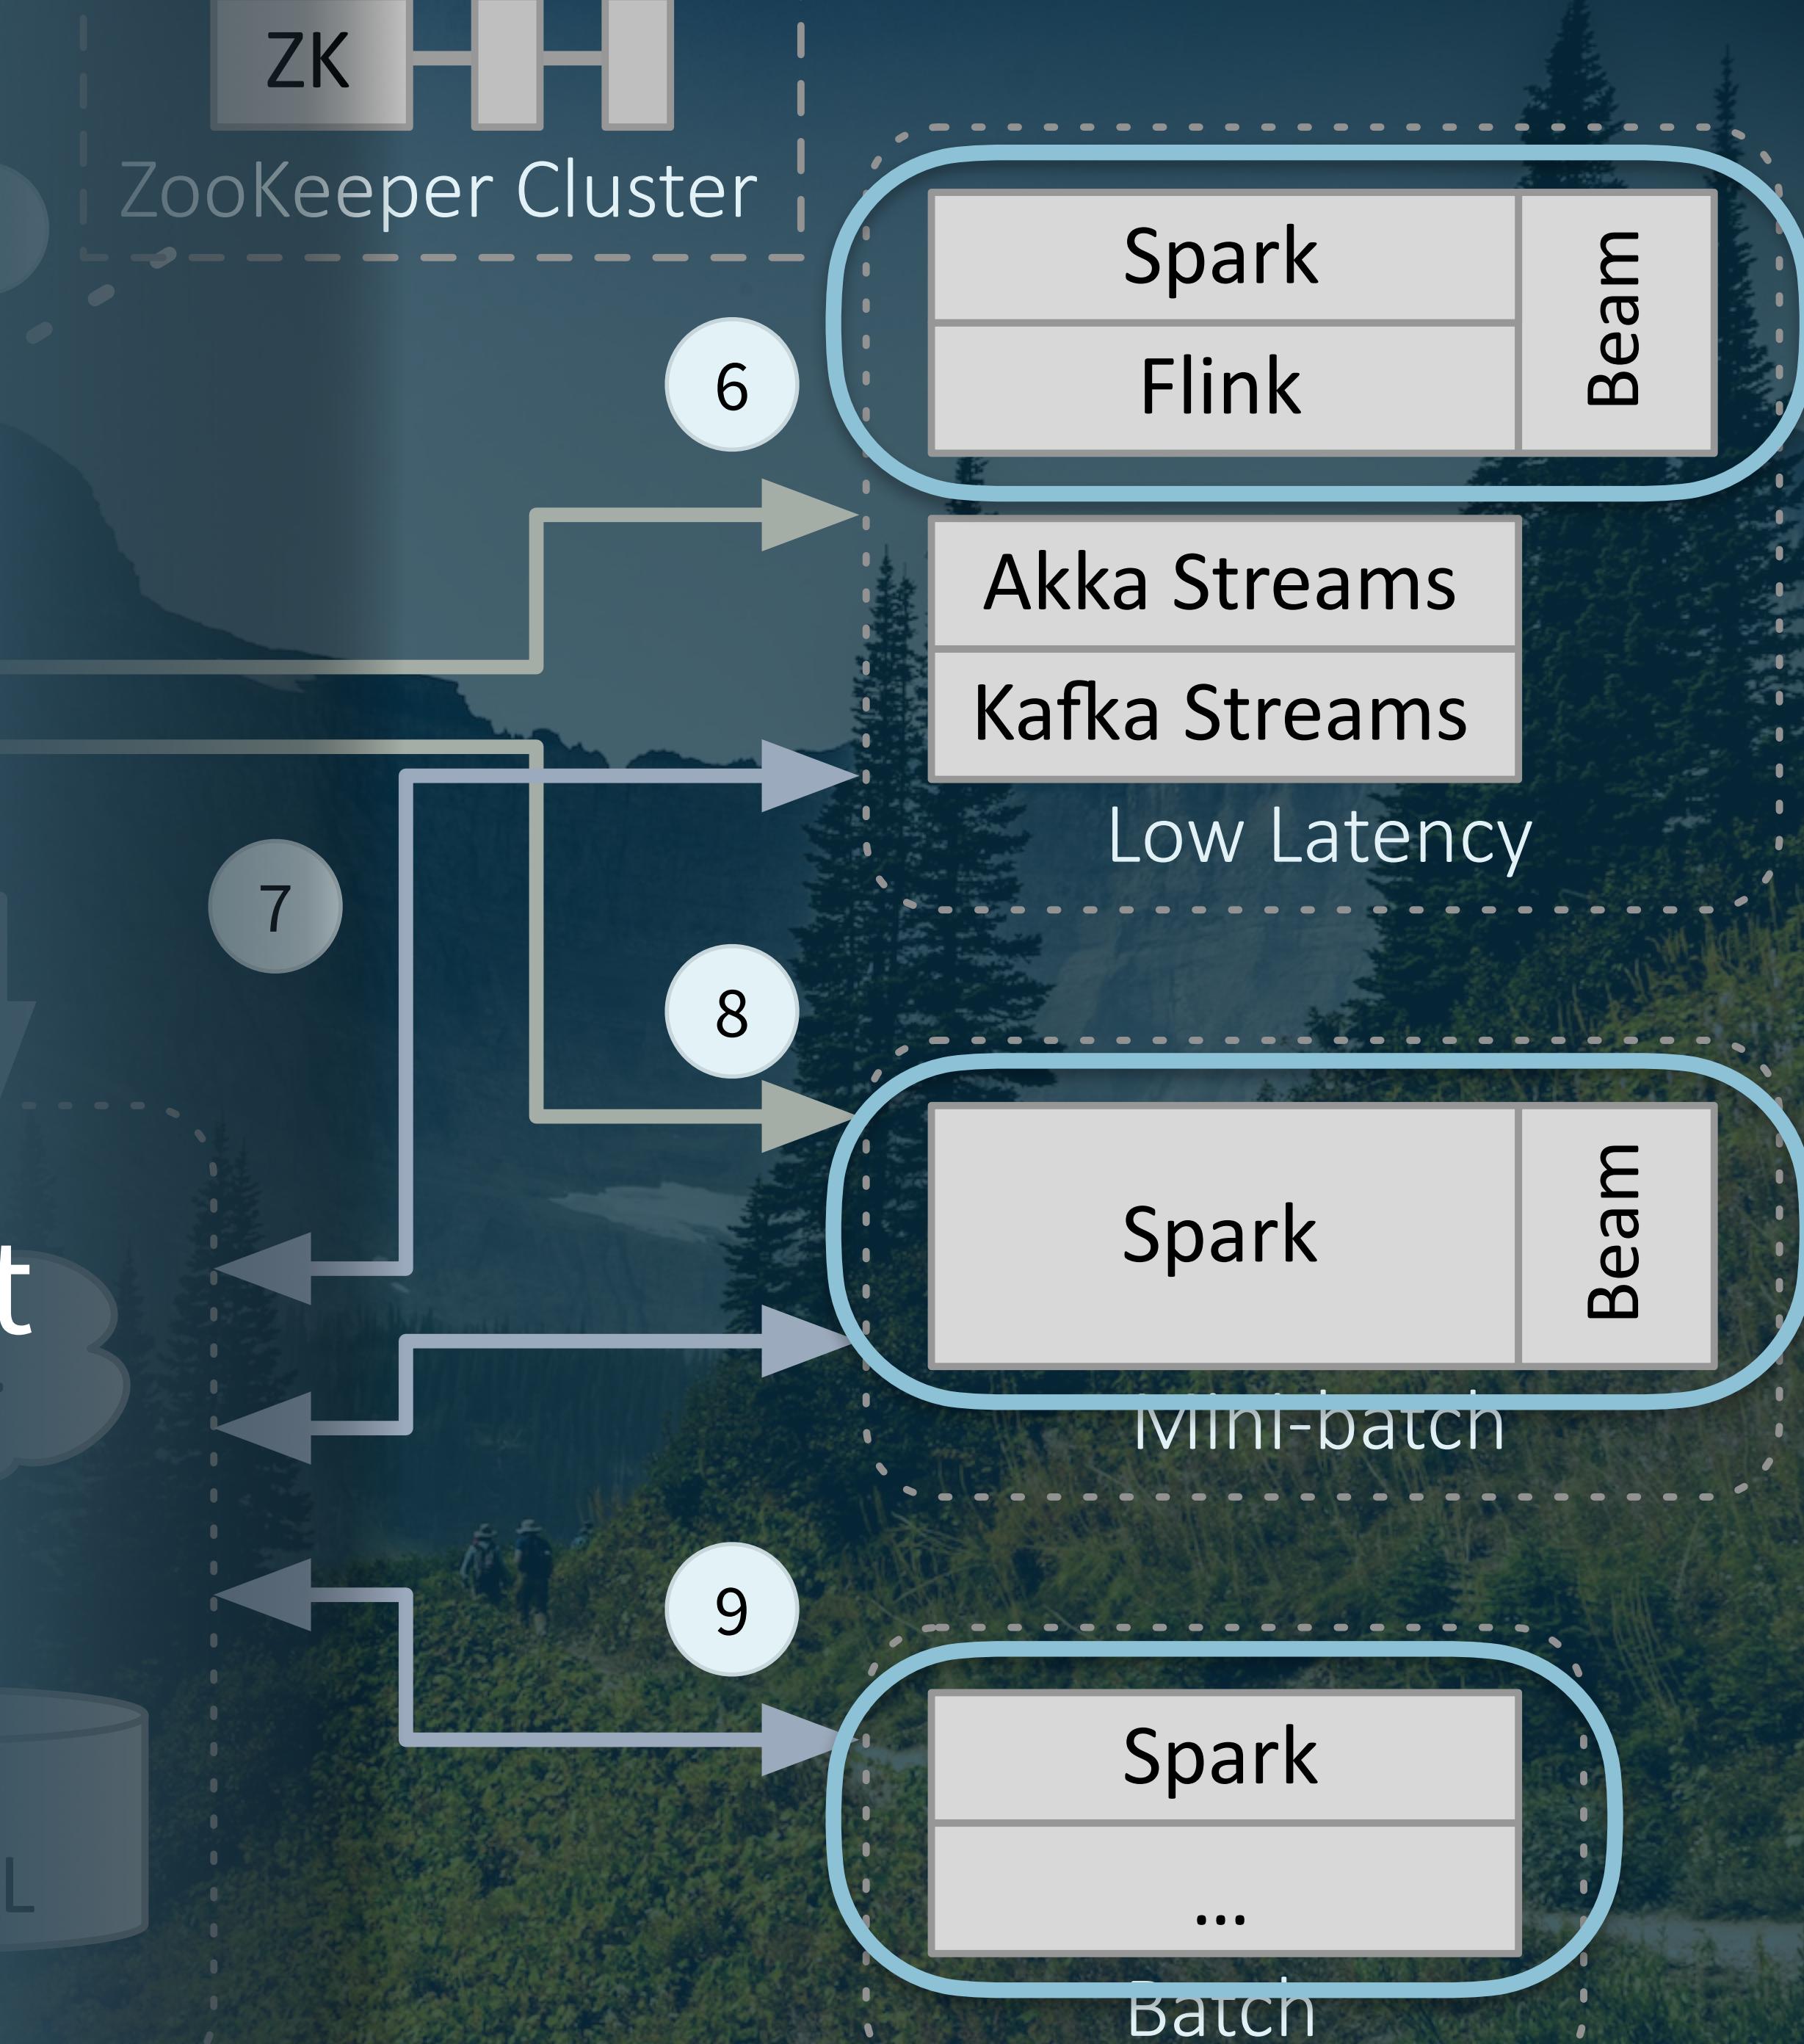


Streaming Engines

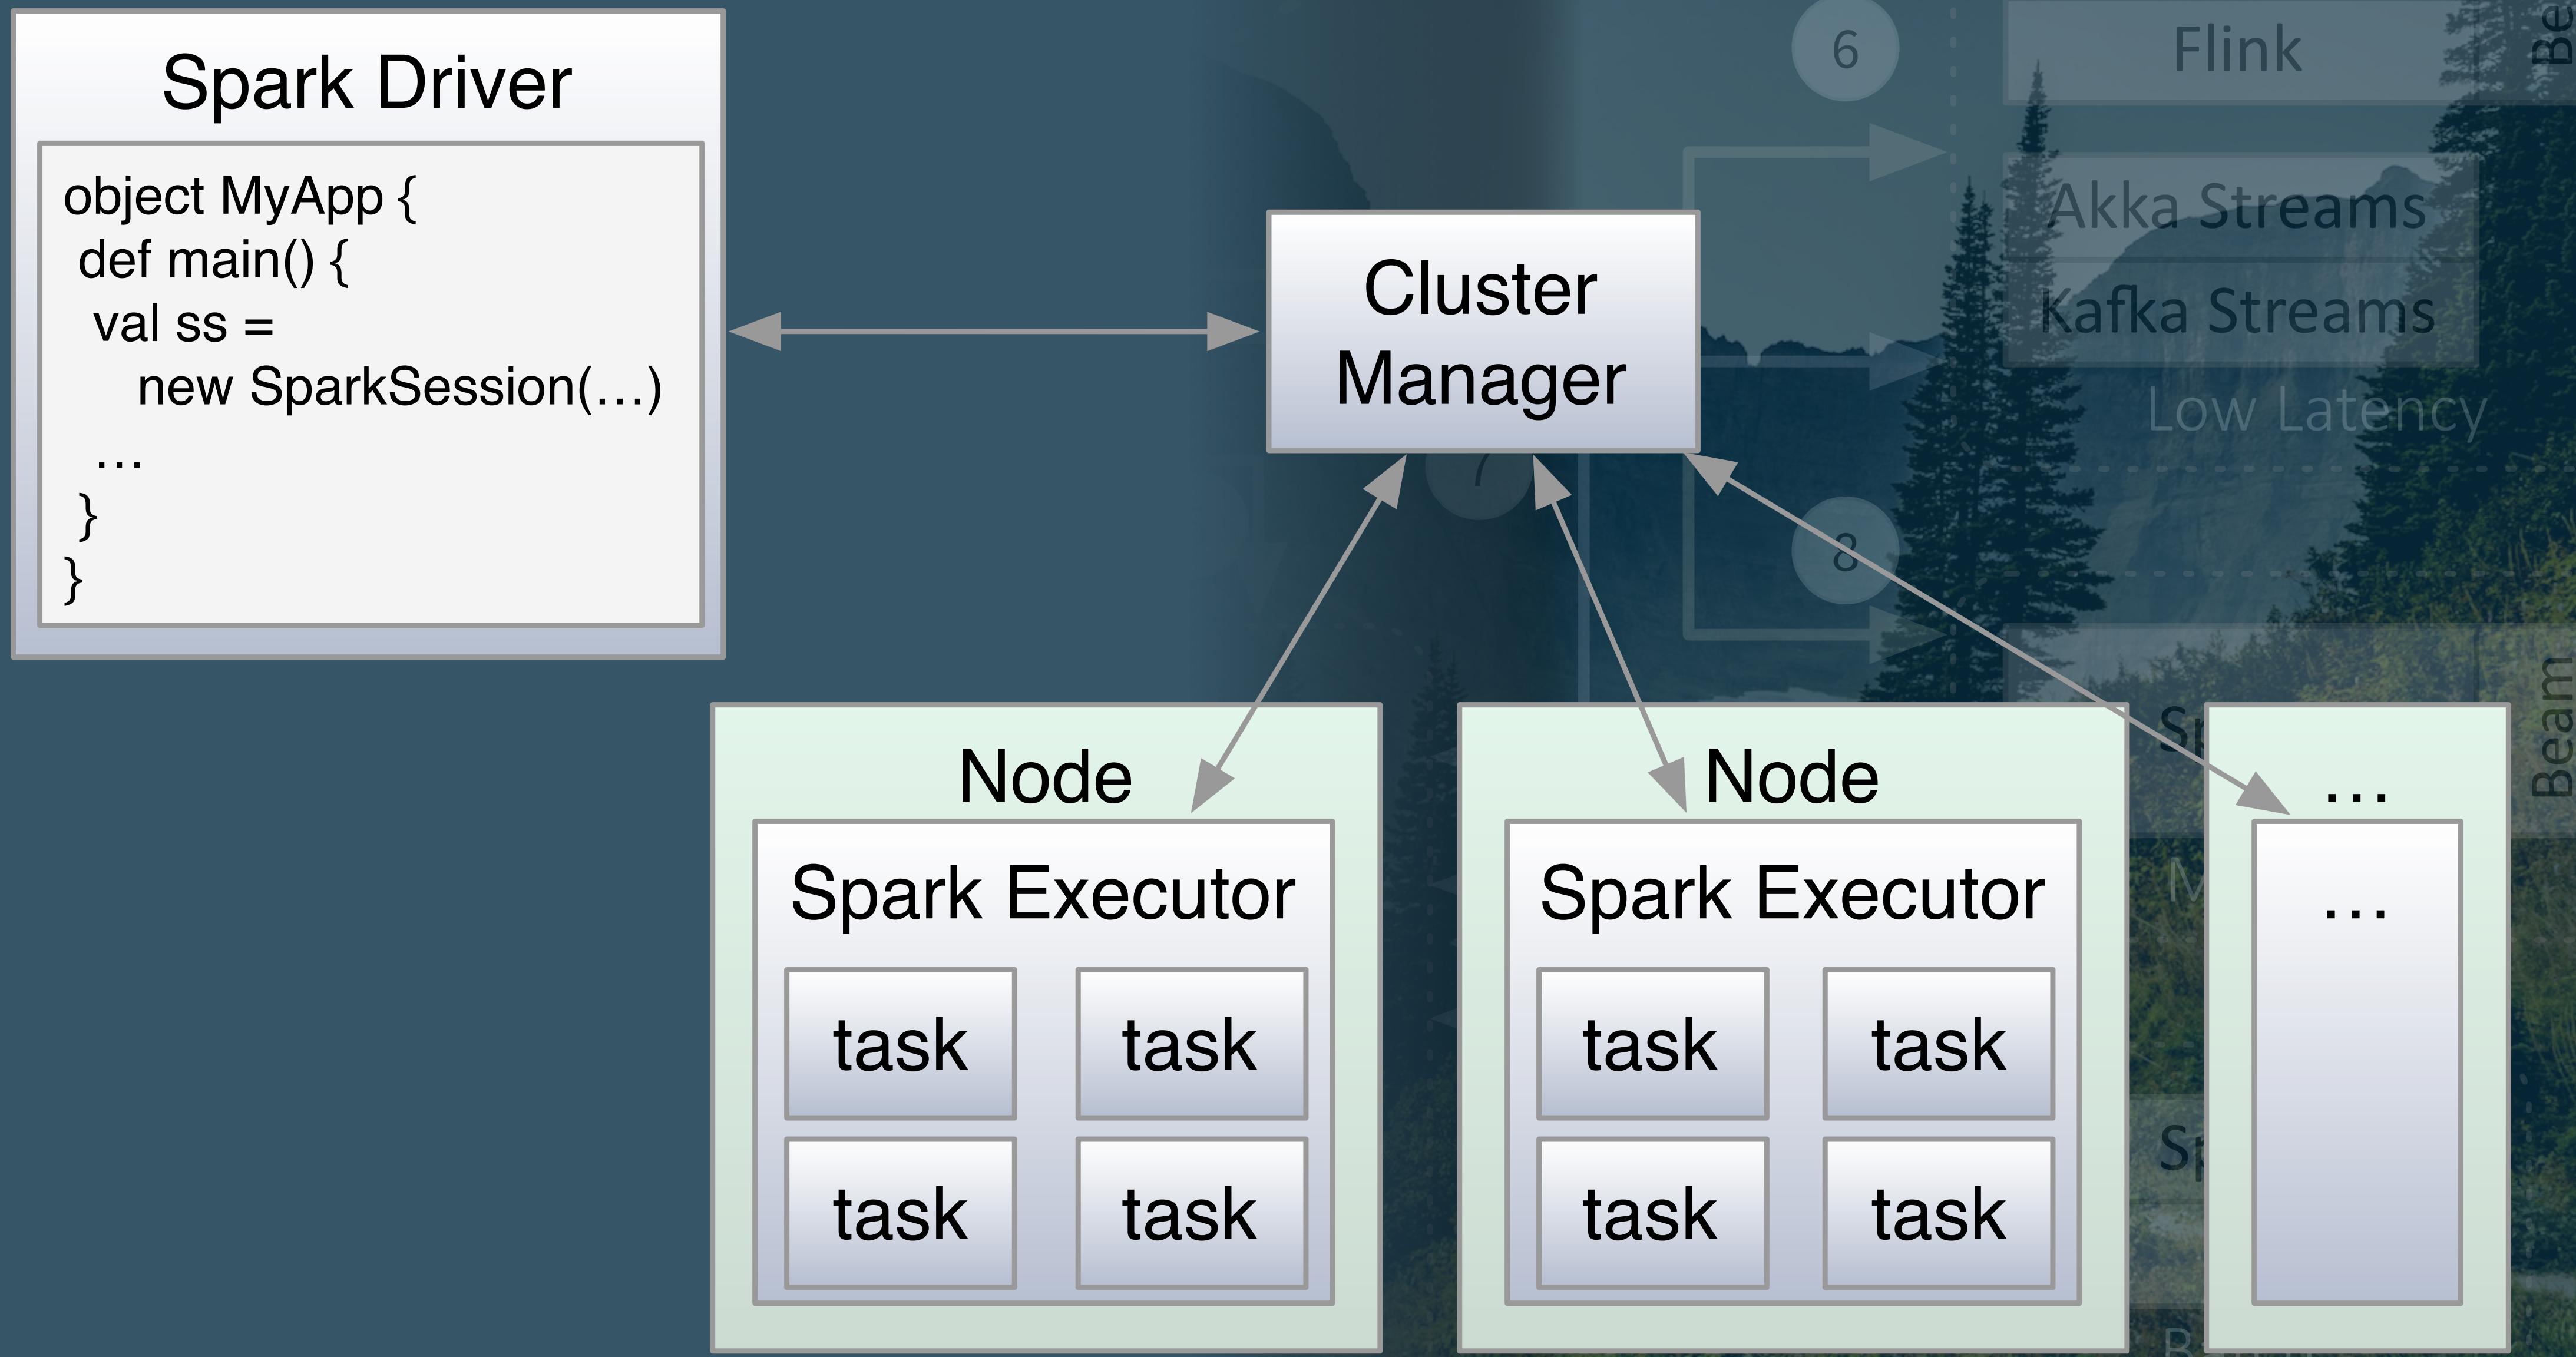


Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.

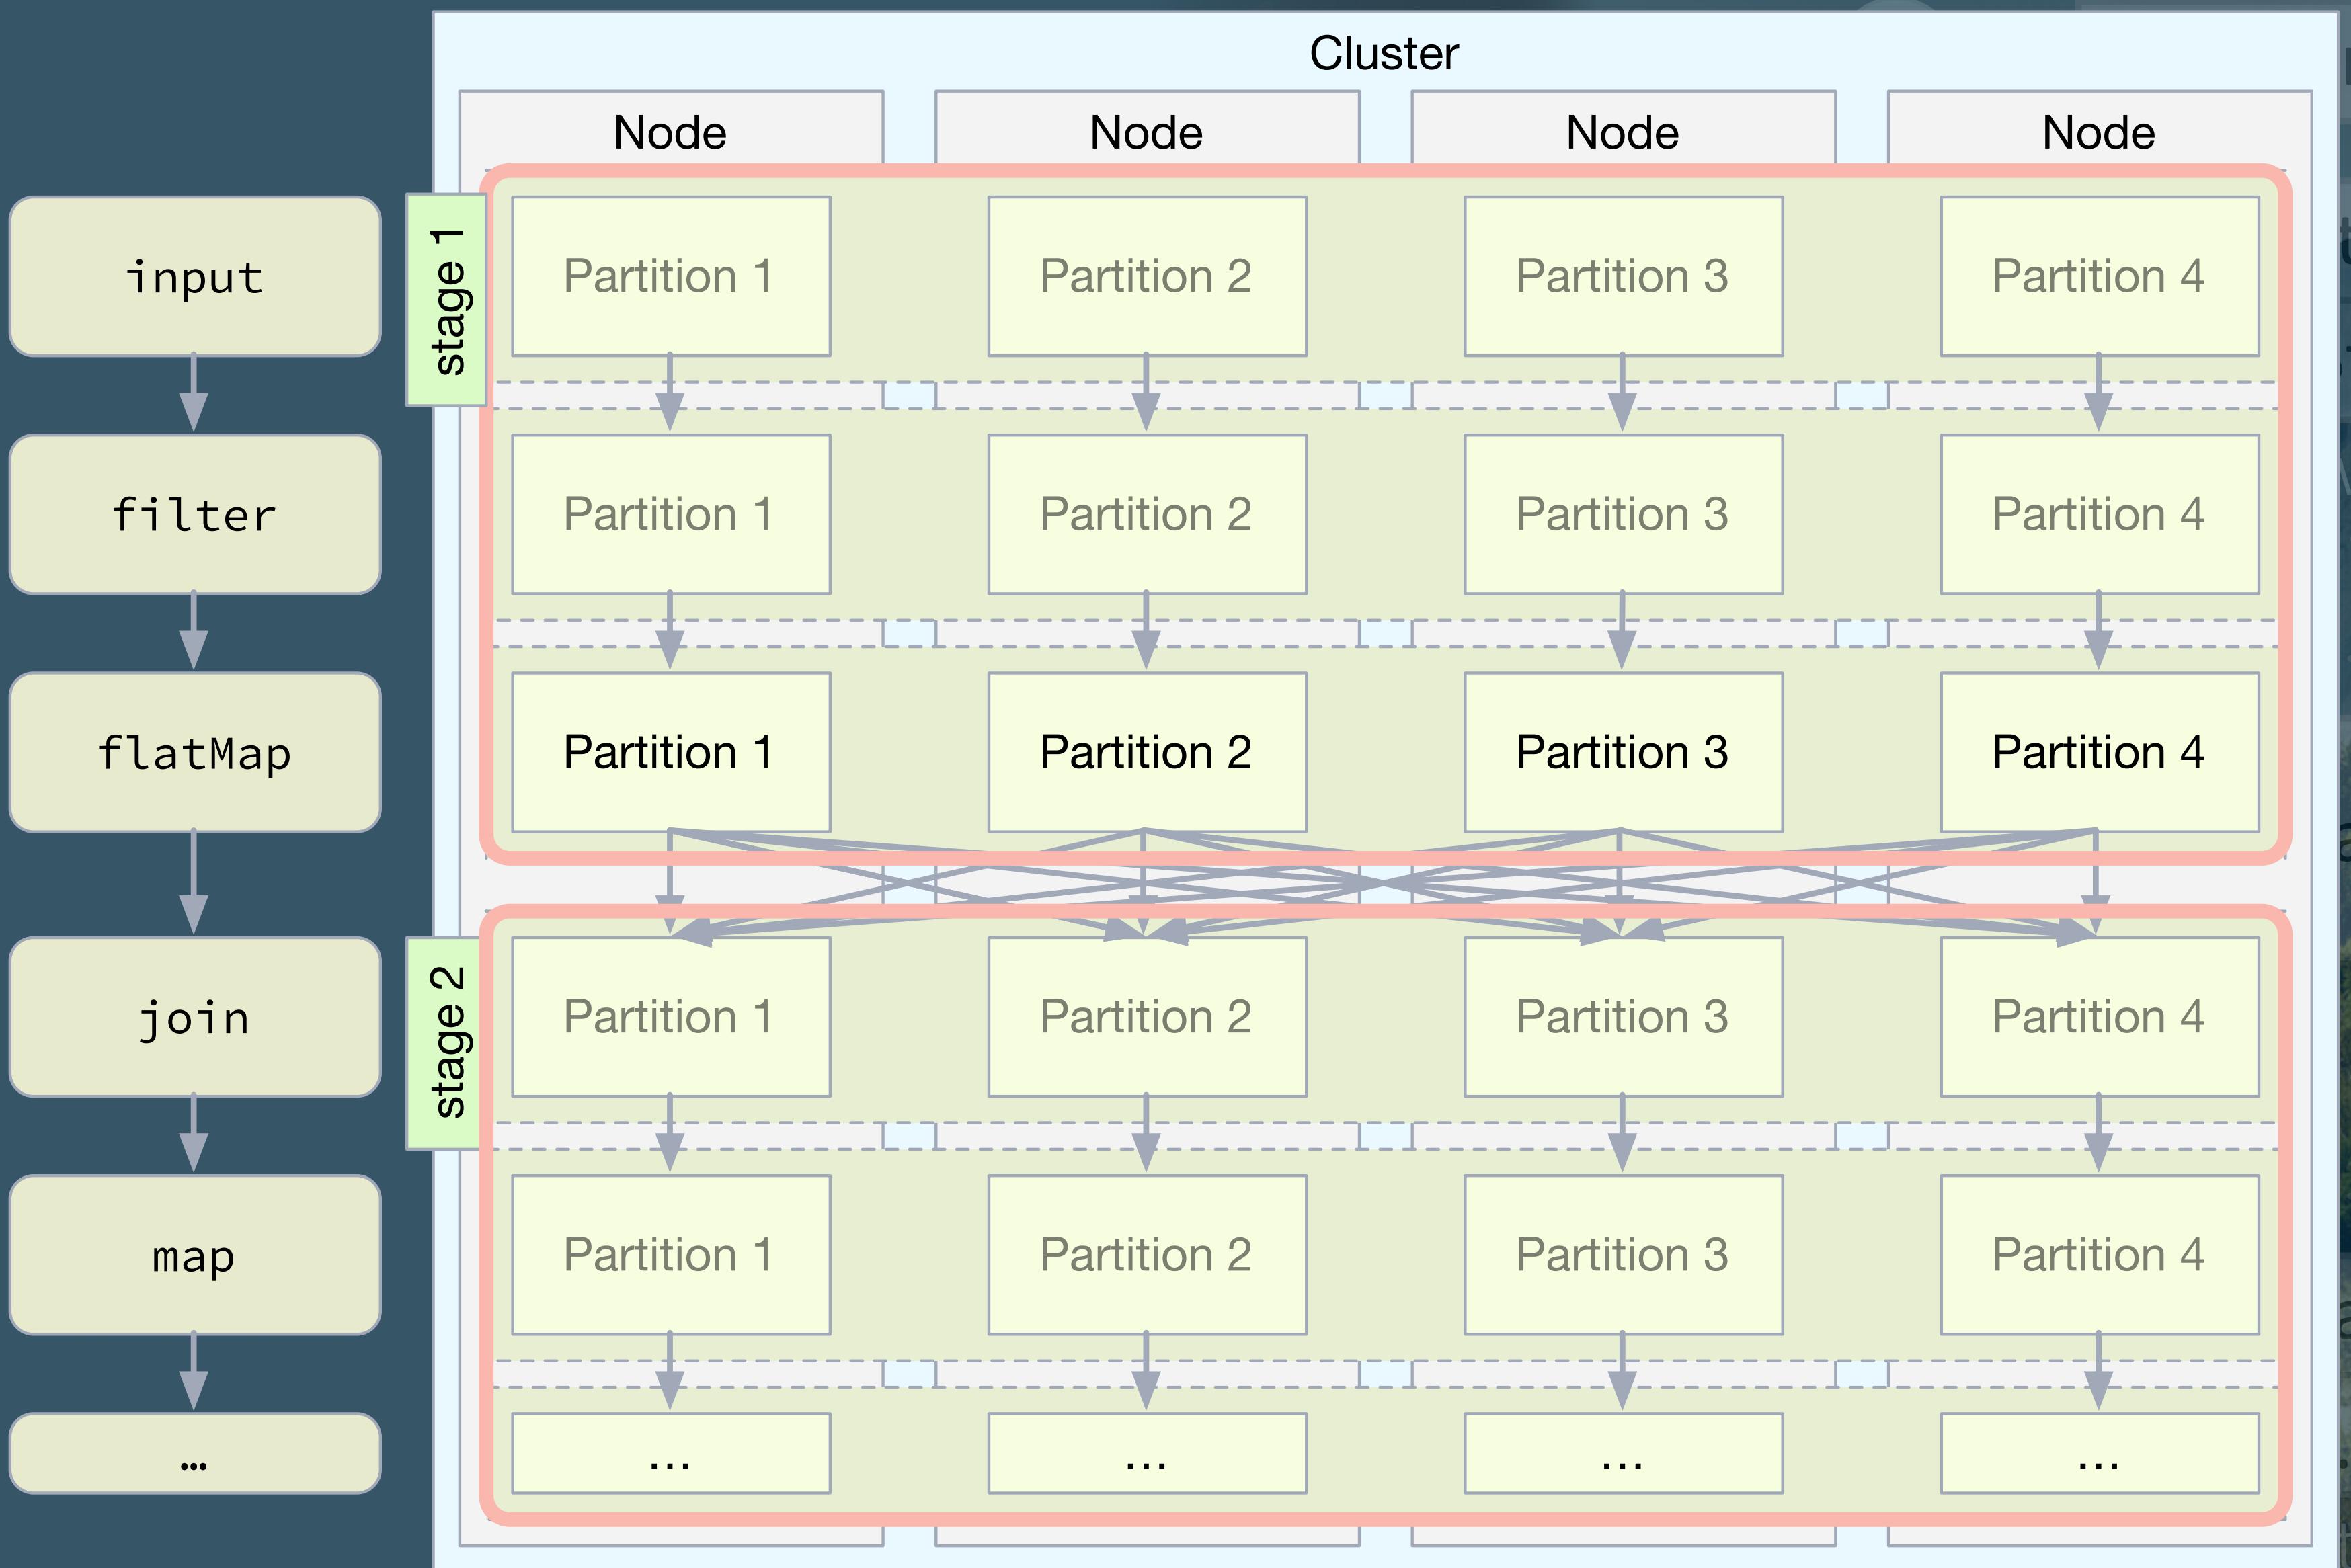
Beam - similar. Google's project that has been instrumental in defining streaming semantics.



They do a *lot* (Spark example)

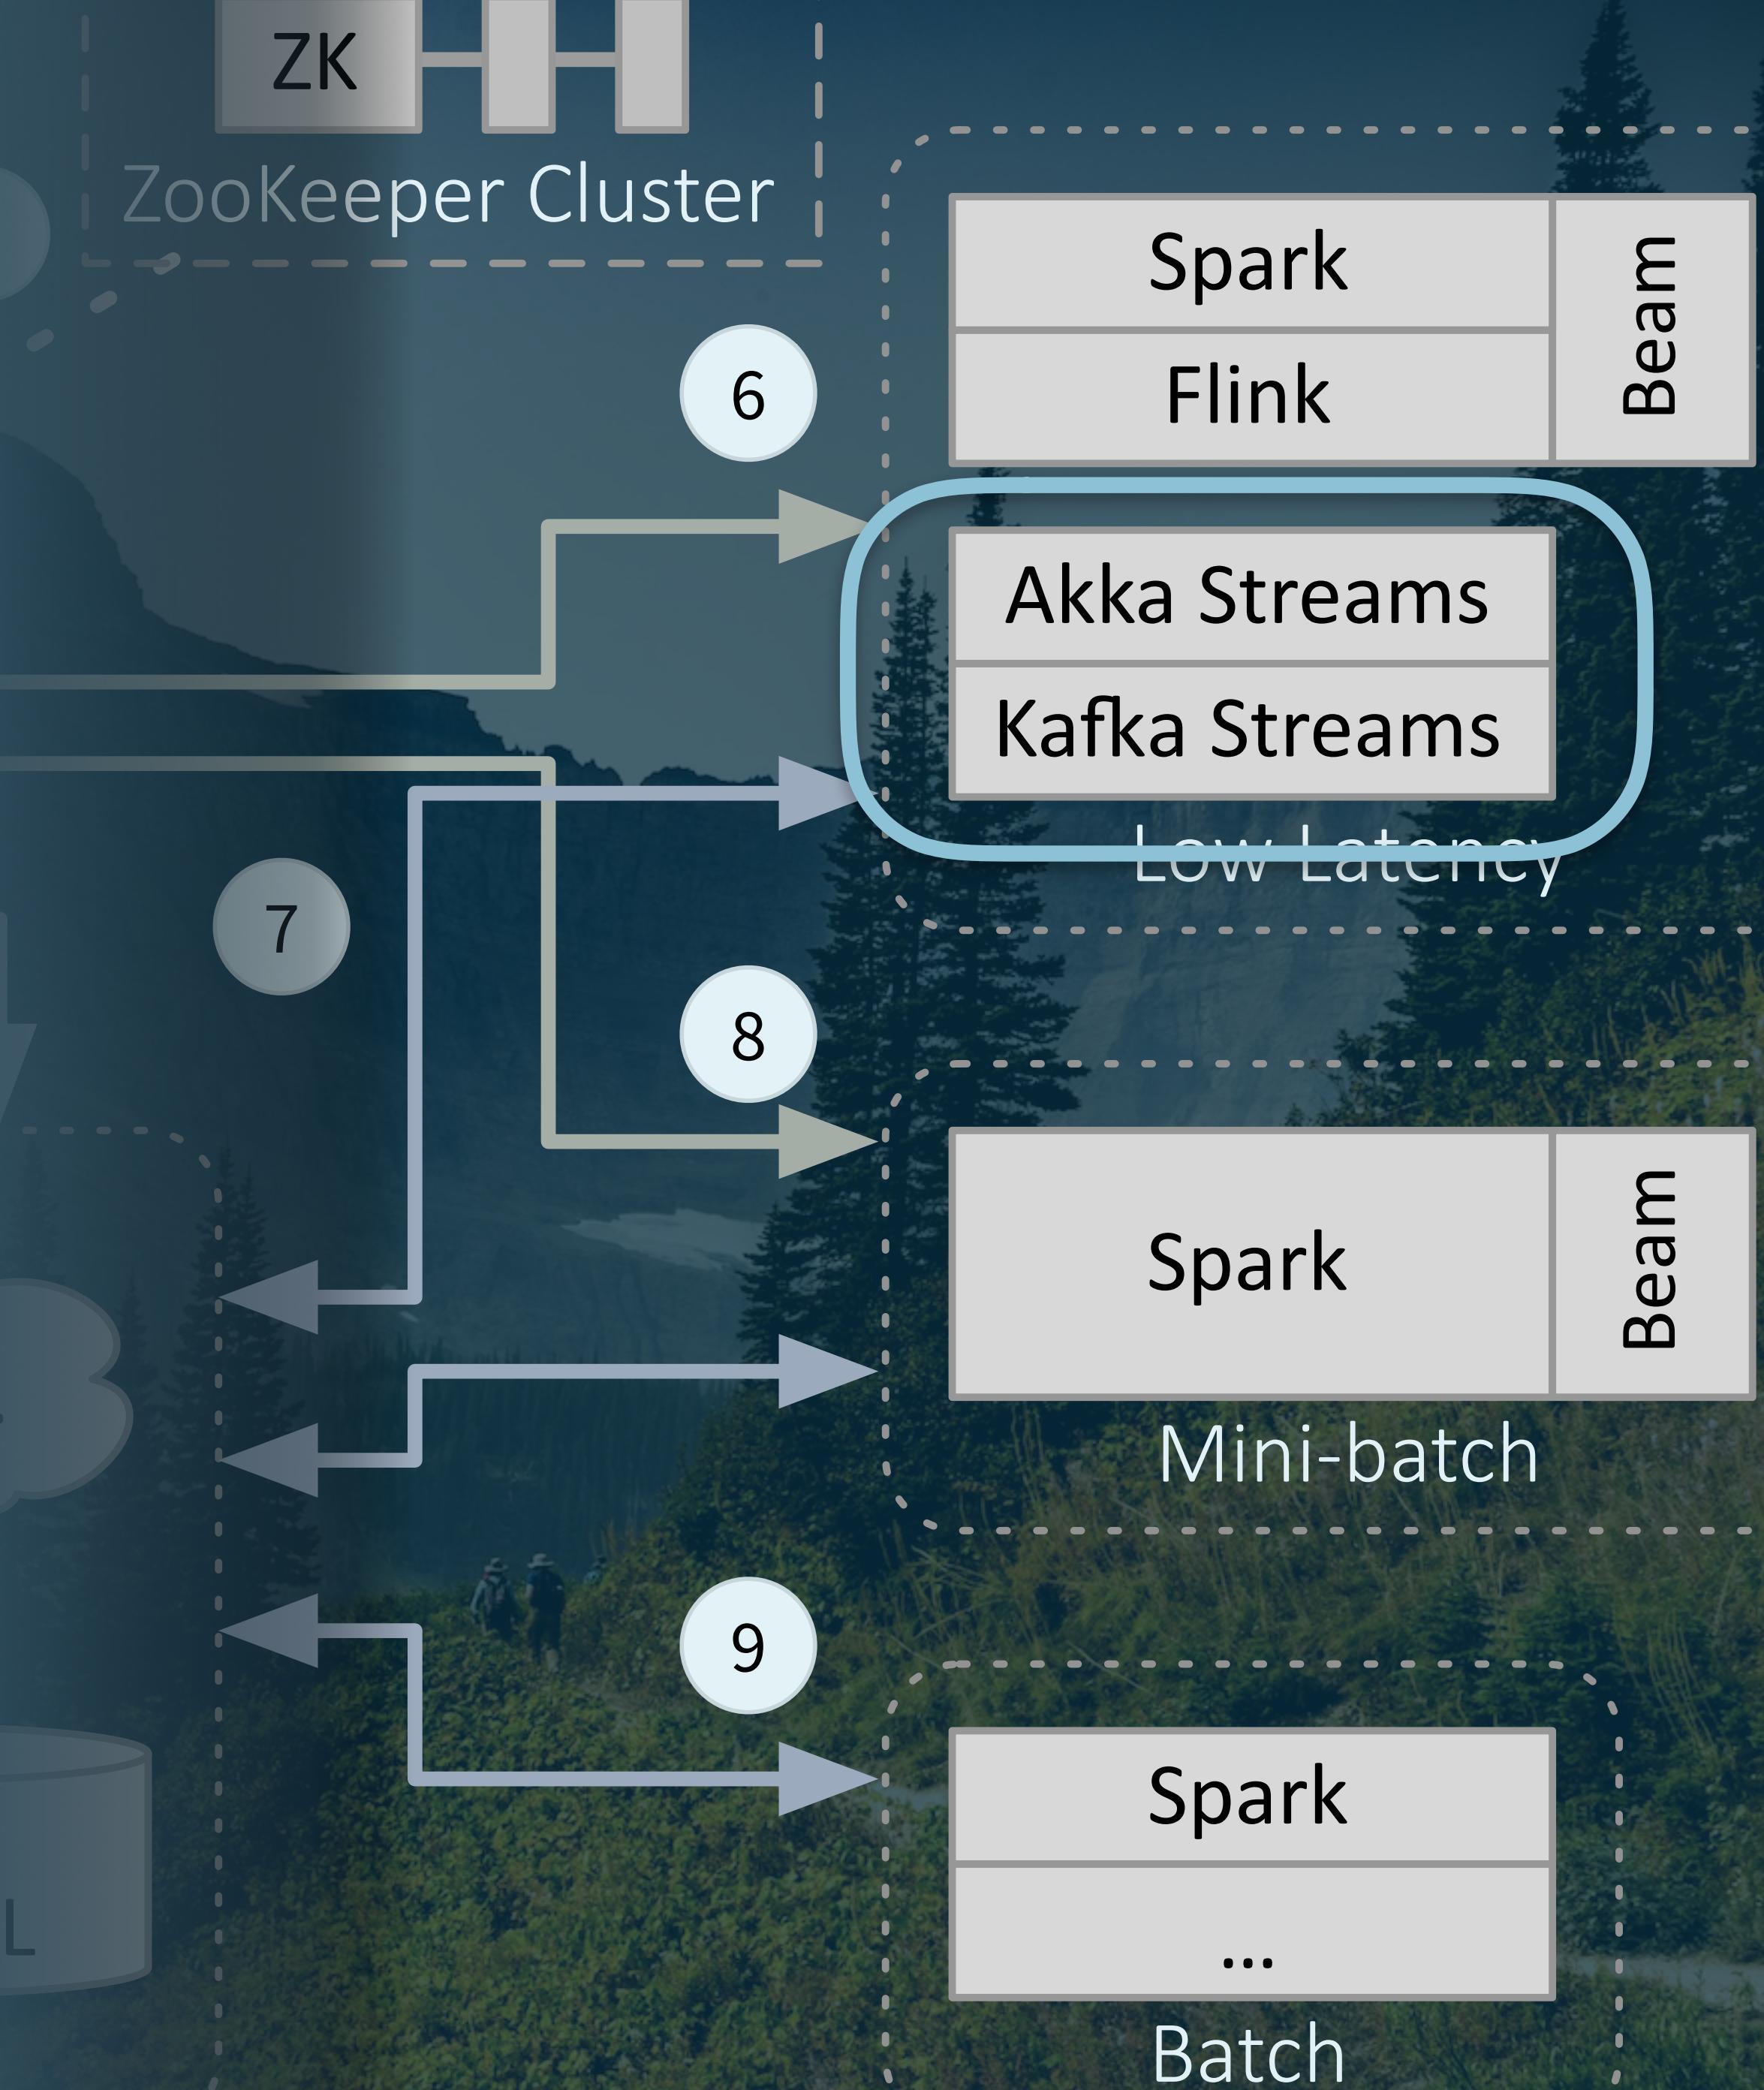


They do a lot (Spark example)



Streaming Engines

Akka Streams, Kafka Streams - libraries for “data-centric microservices”. Smaller scale, but great flexibility



Microservice All the Things!!



Scott Hanselman

@shanselman

Follow

Microservices, for when your in-process methods have too little latency.

Dave Cheney @davecheney

Microservices, for when function calls are too reliable.

4:11 AM - 25 Feb 2018

207 Retweets **566** Likes



25

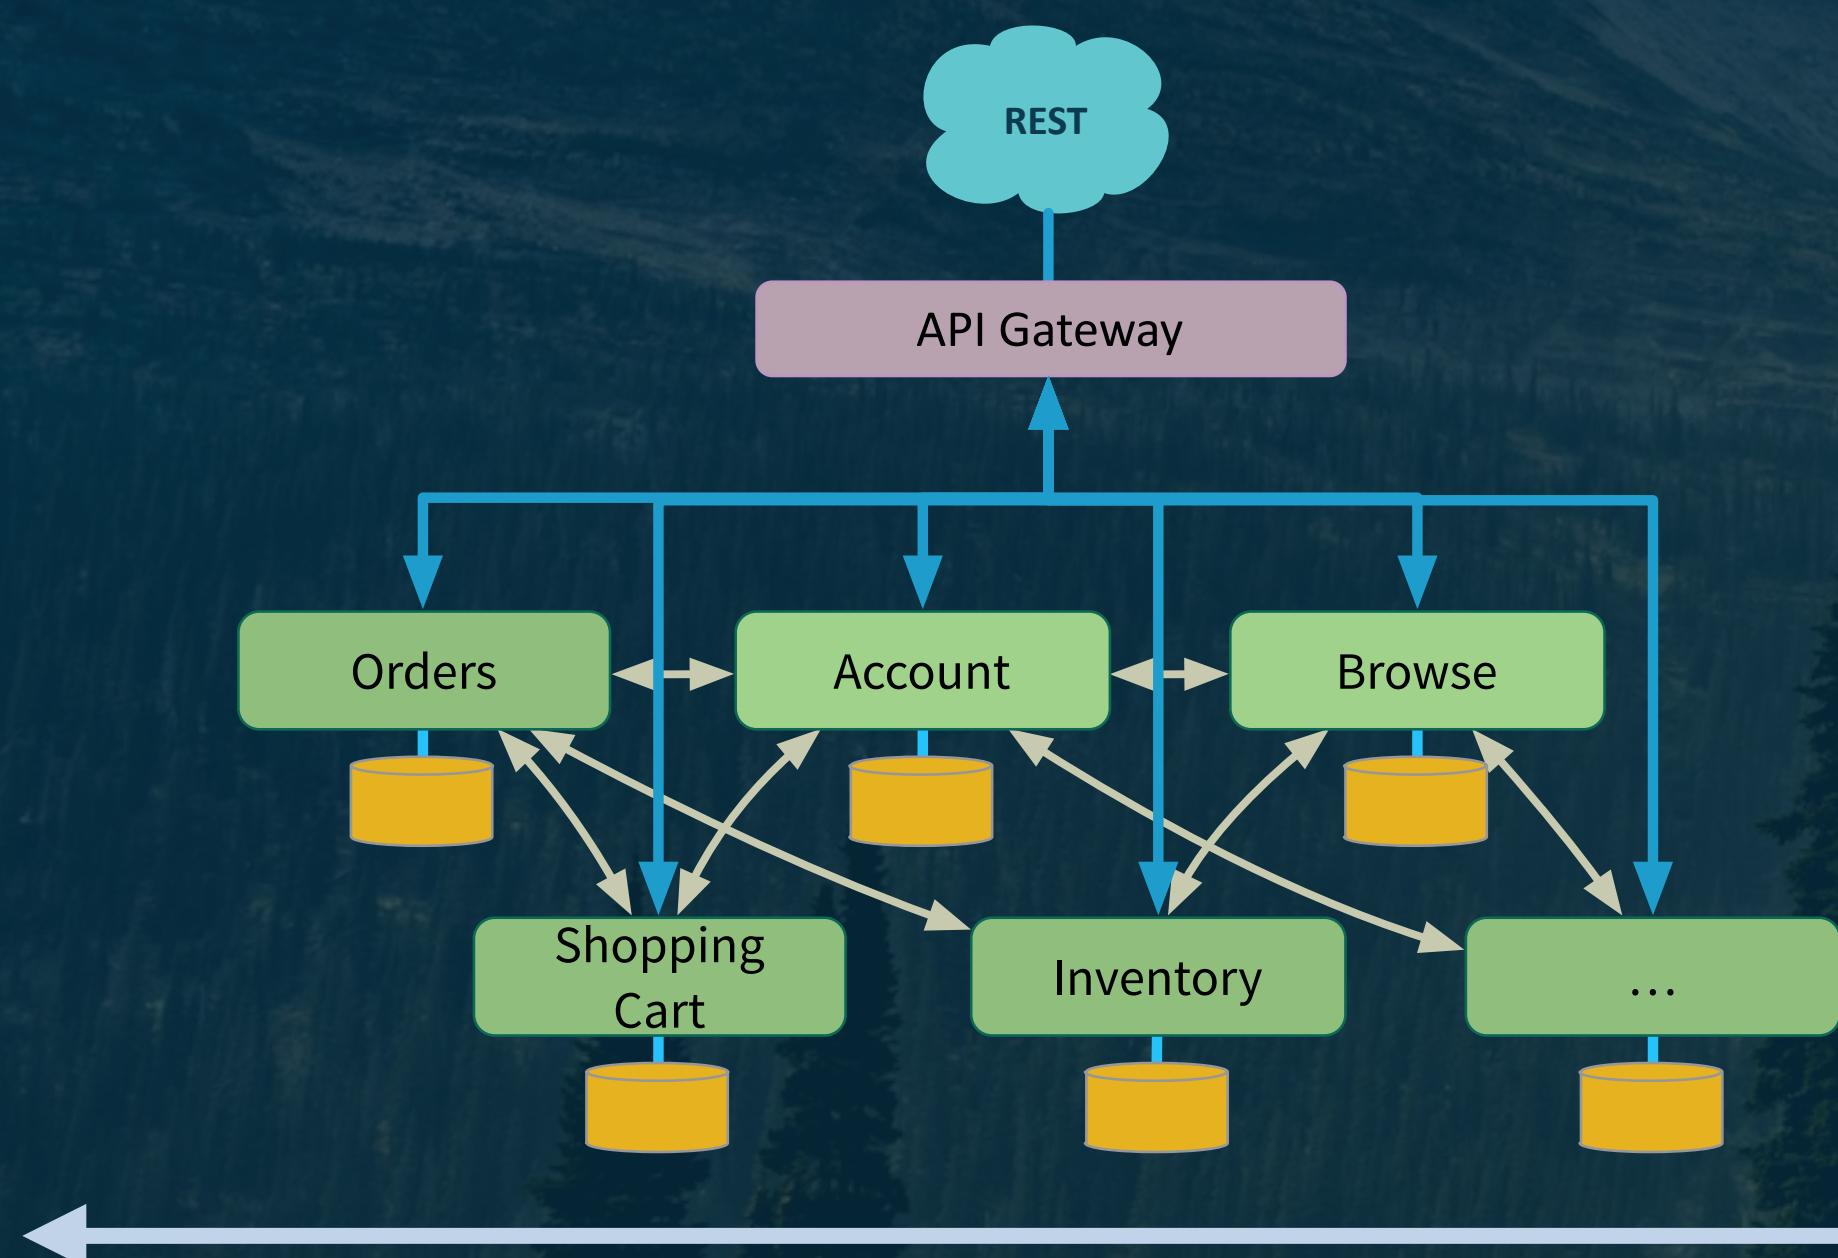
207

566

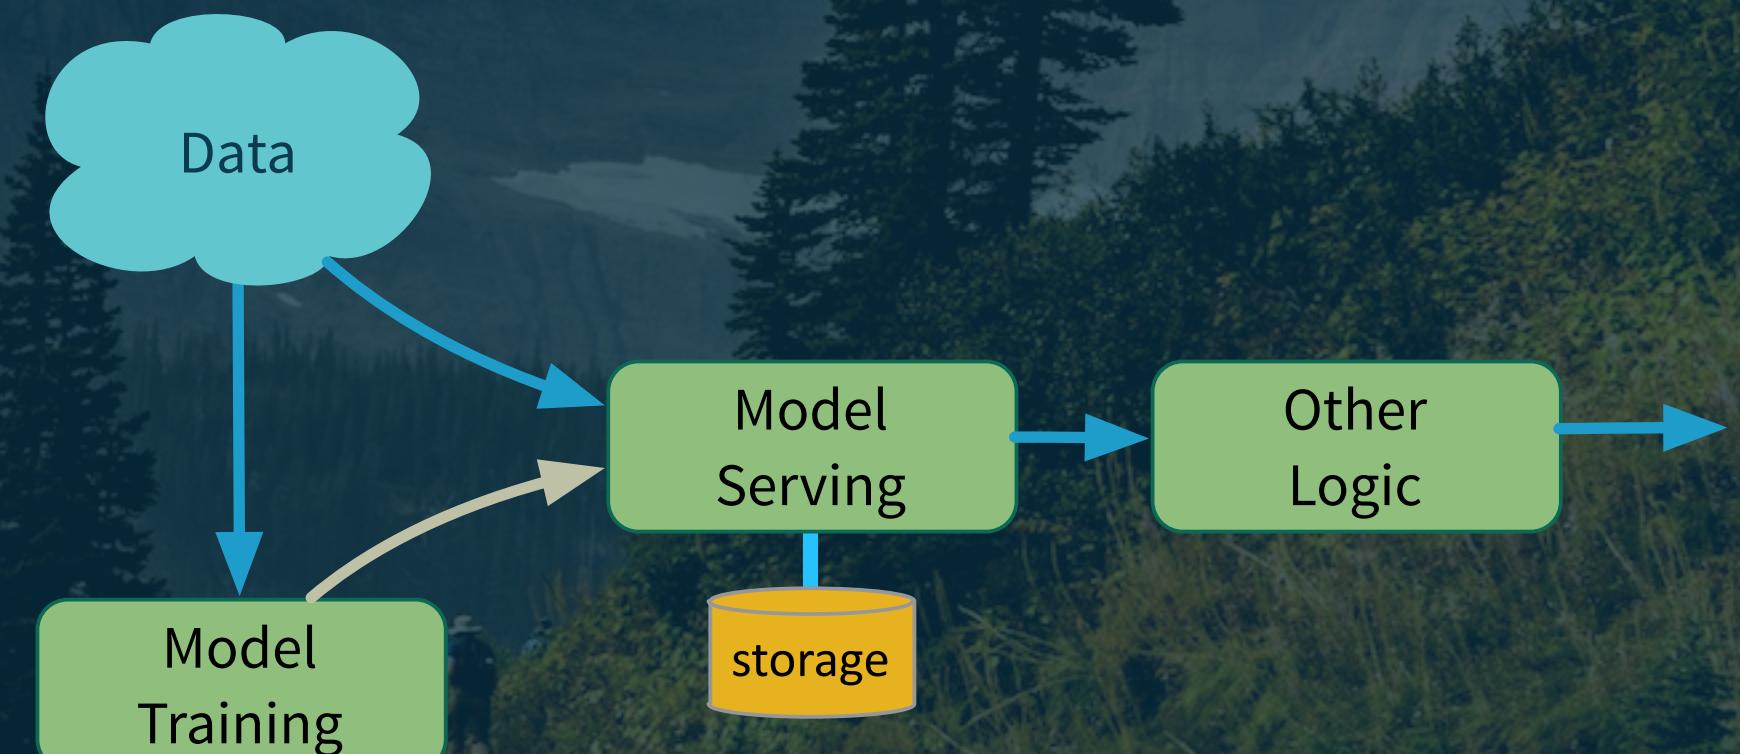


A Spectrum of Microservices

Event-driven μ-services



“Record-centric” μ-services

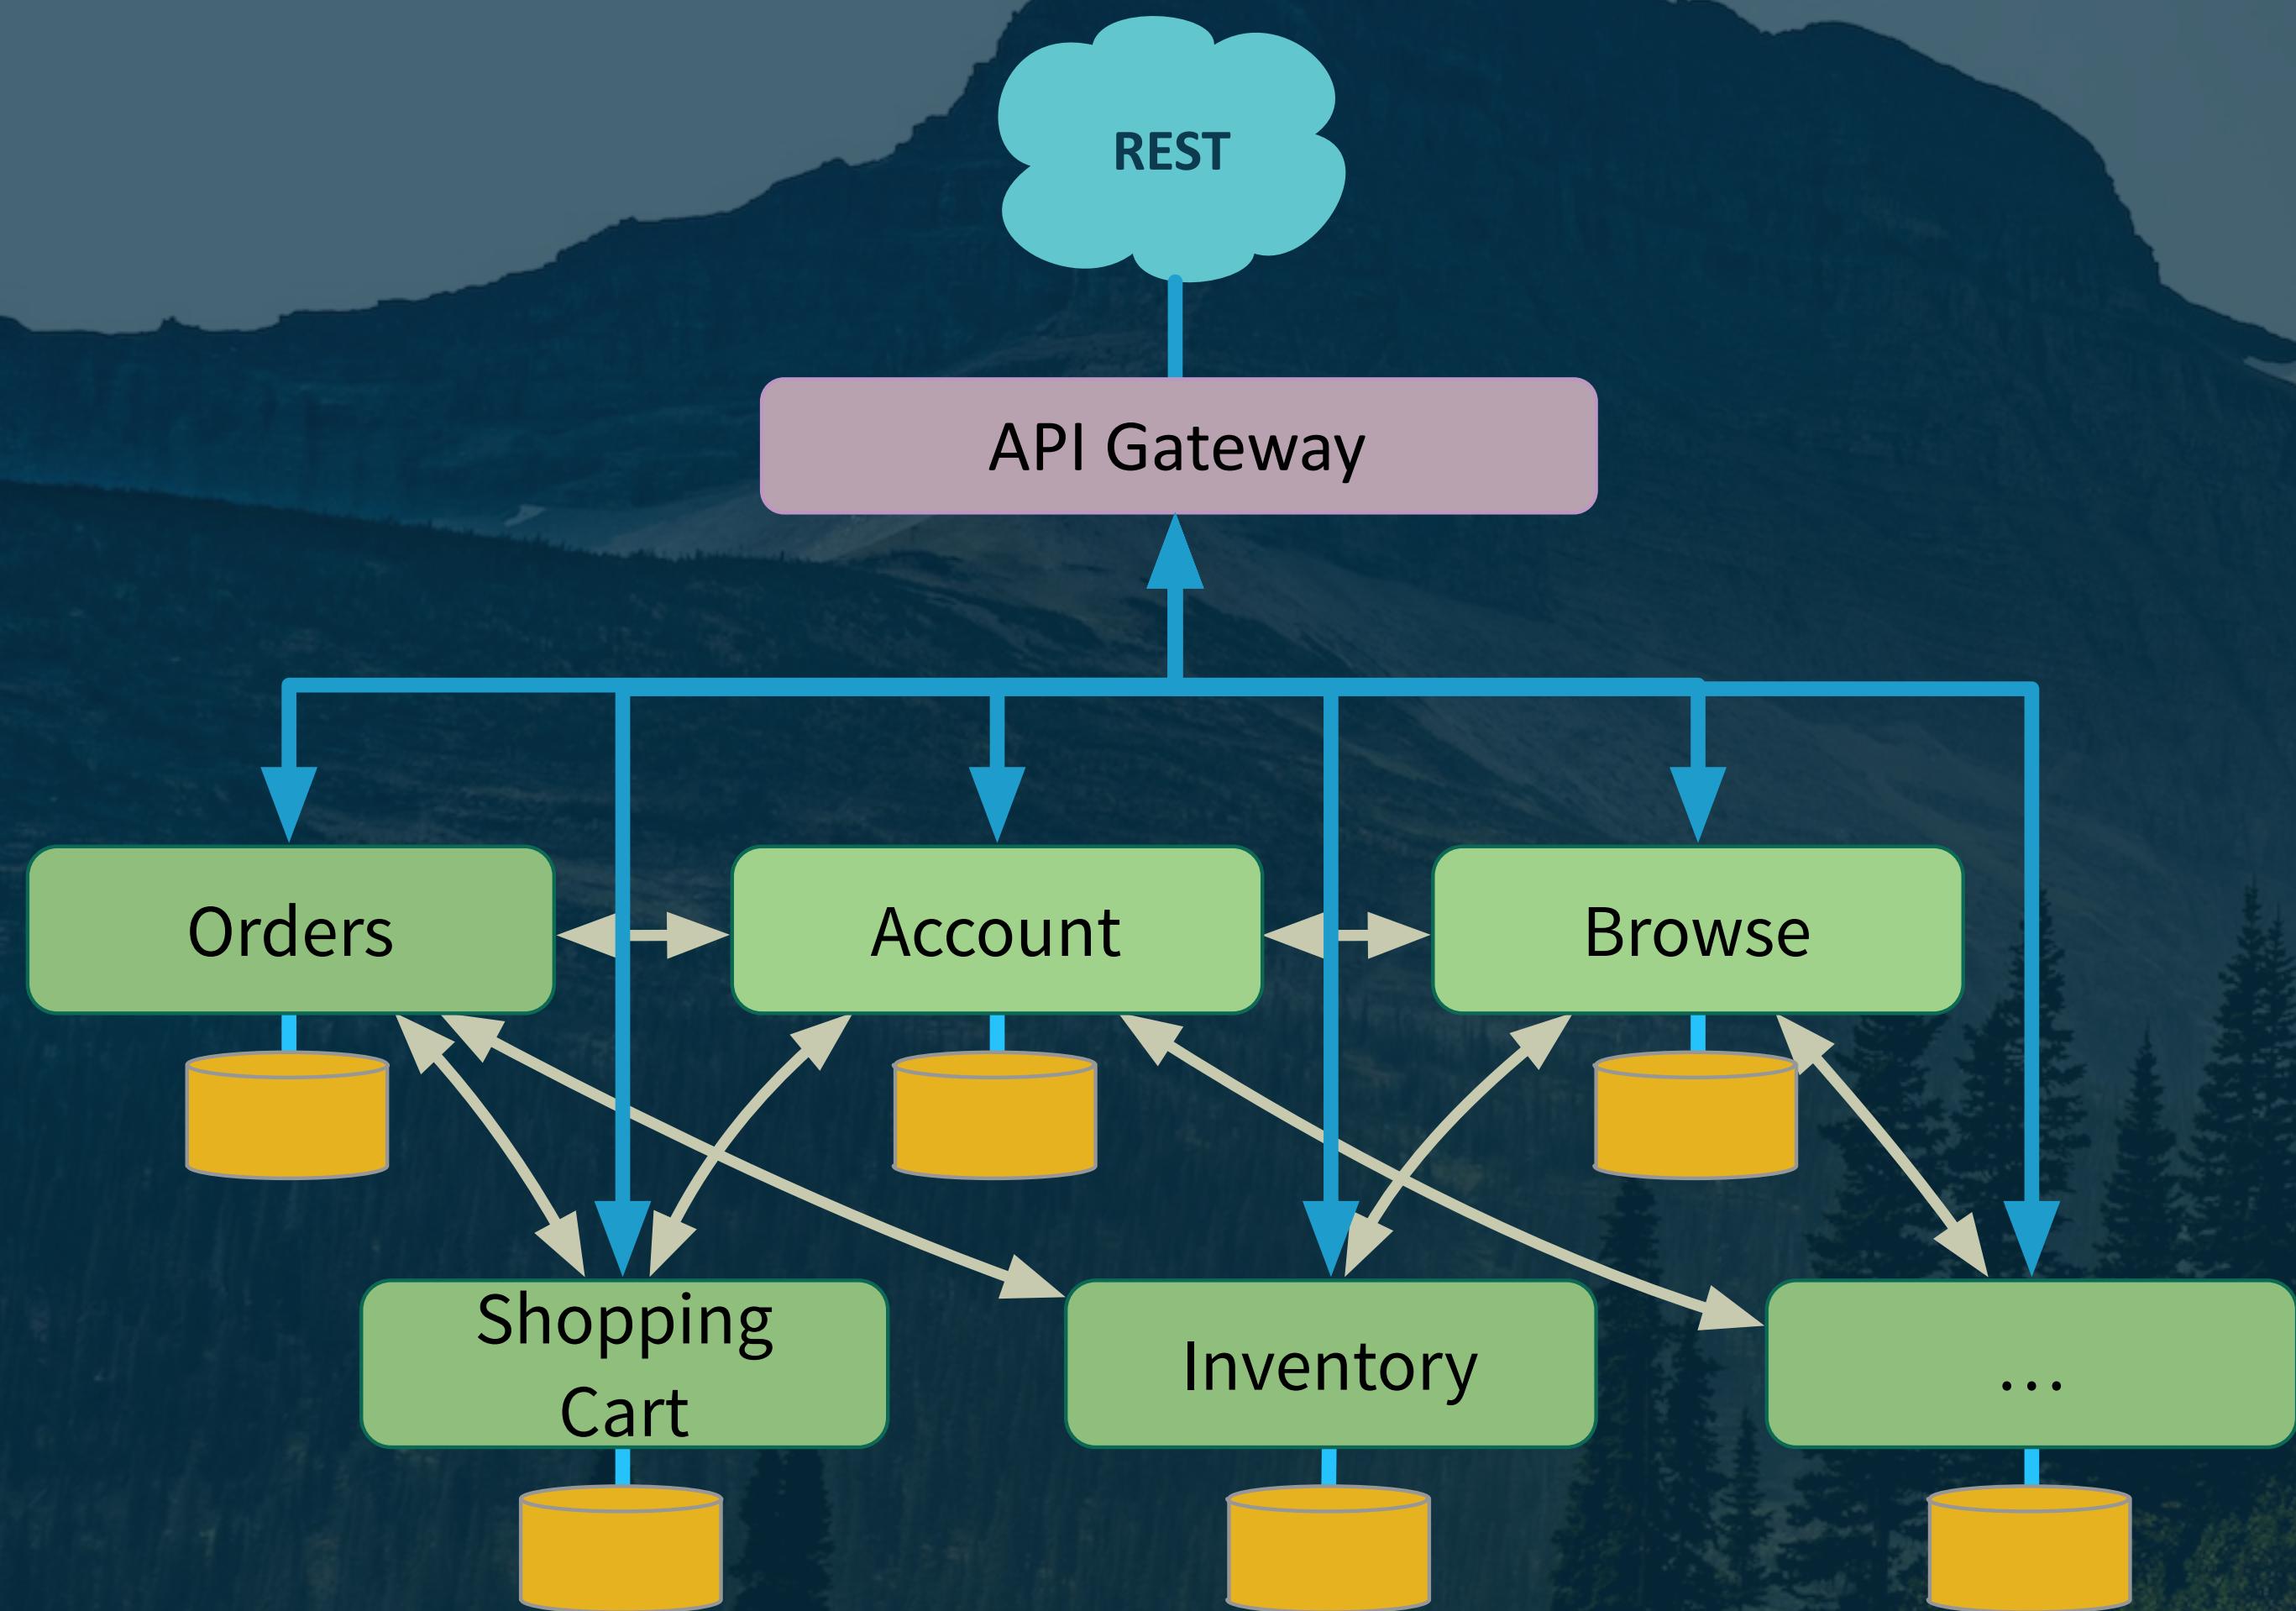


Events

← Data Spectrum →

Records

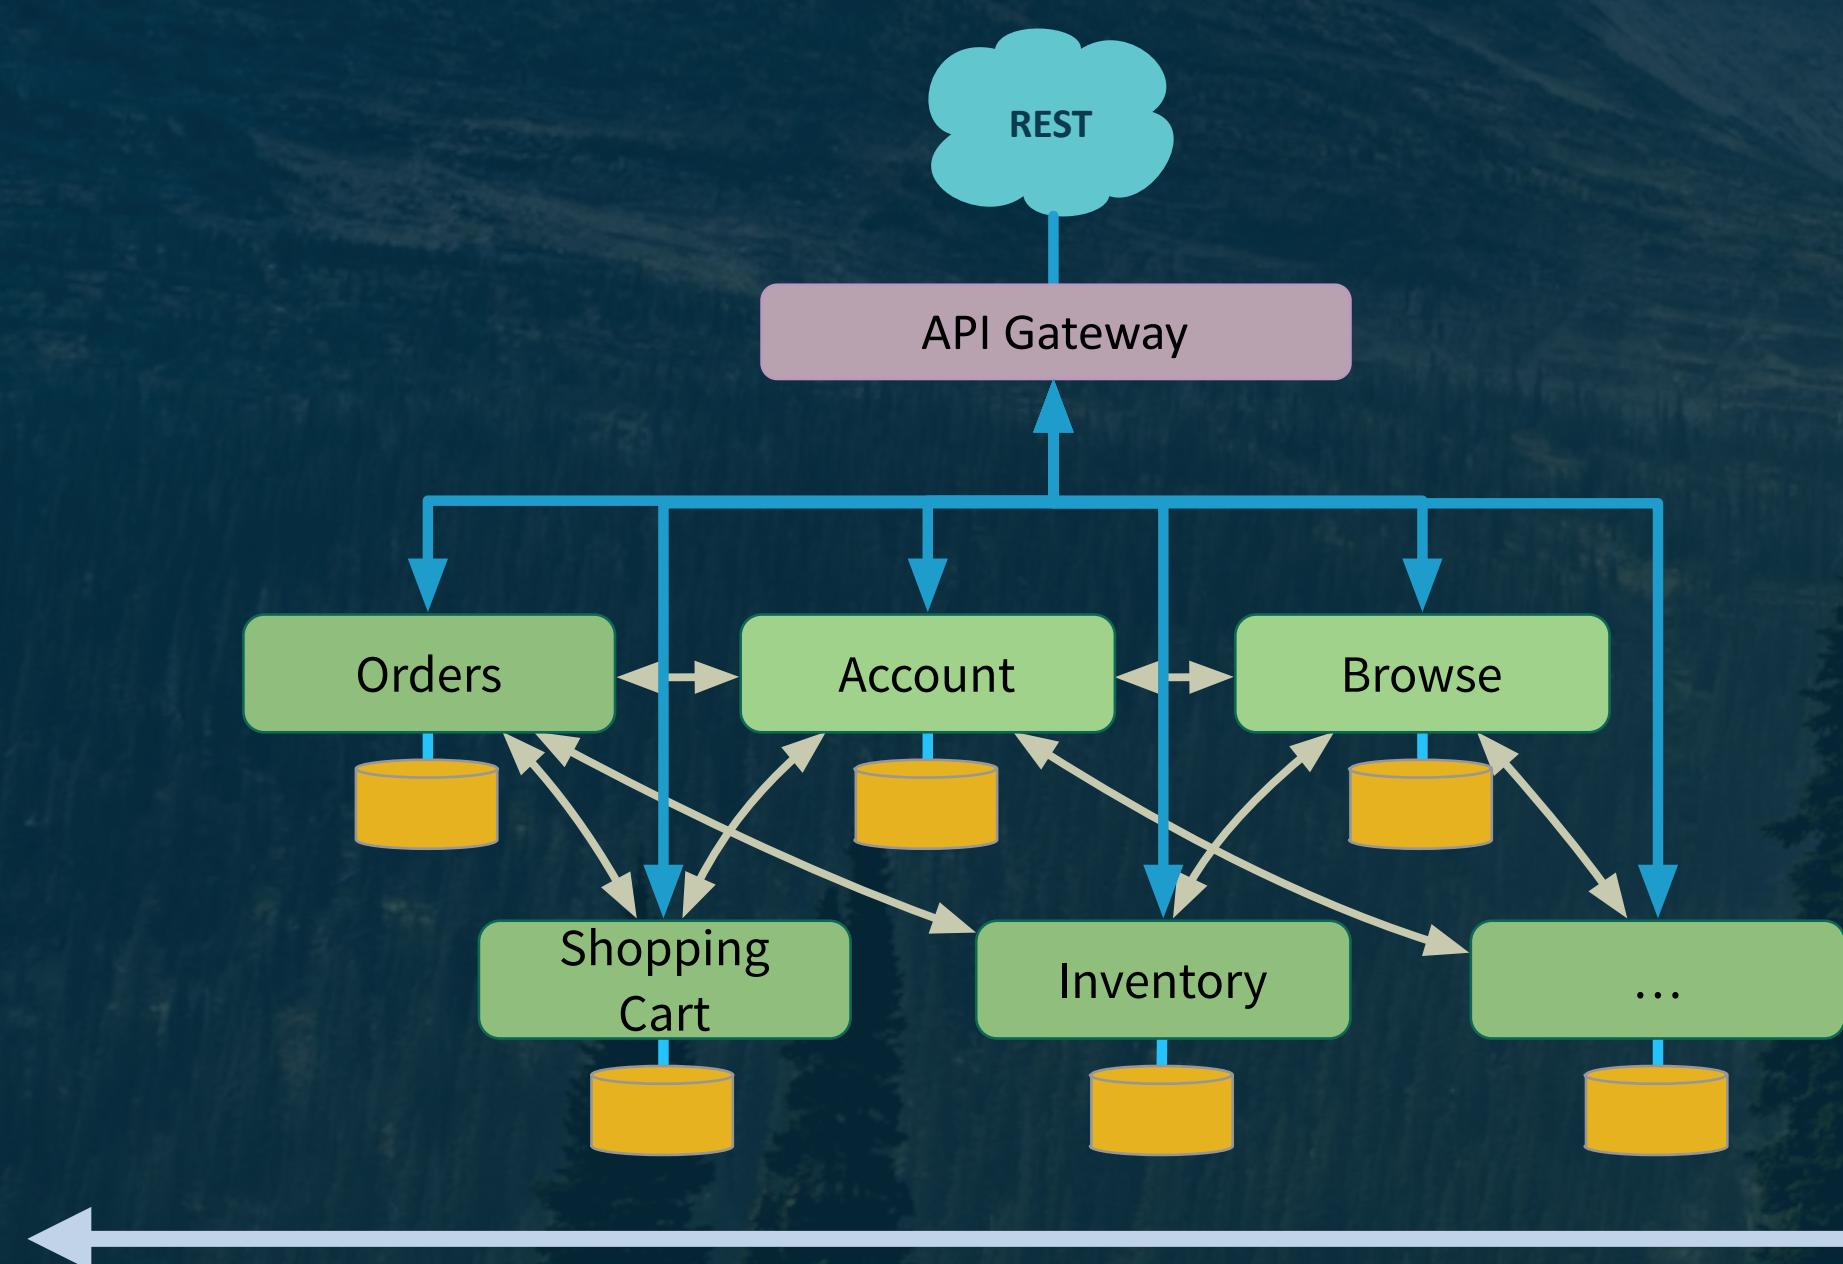
A Spectrum of Microservices



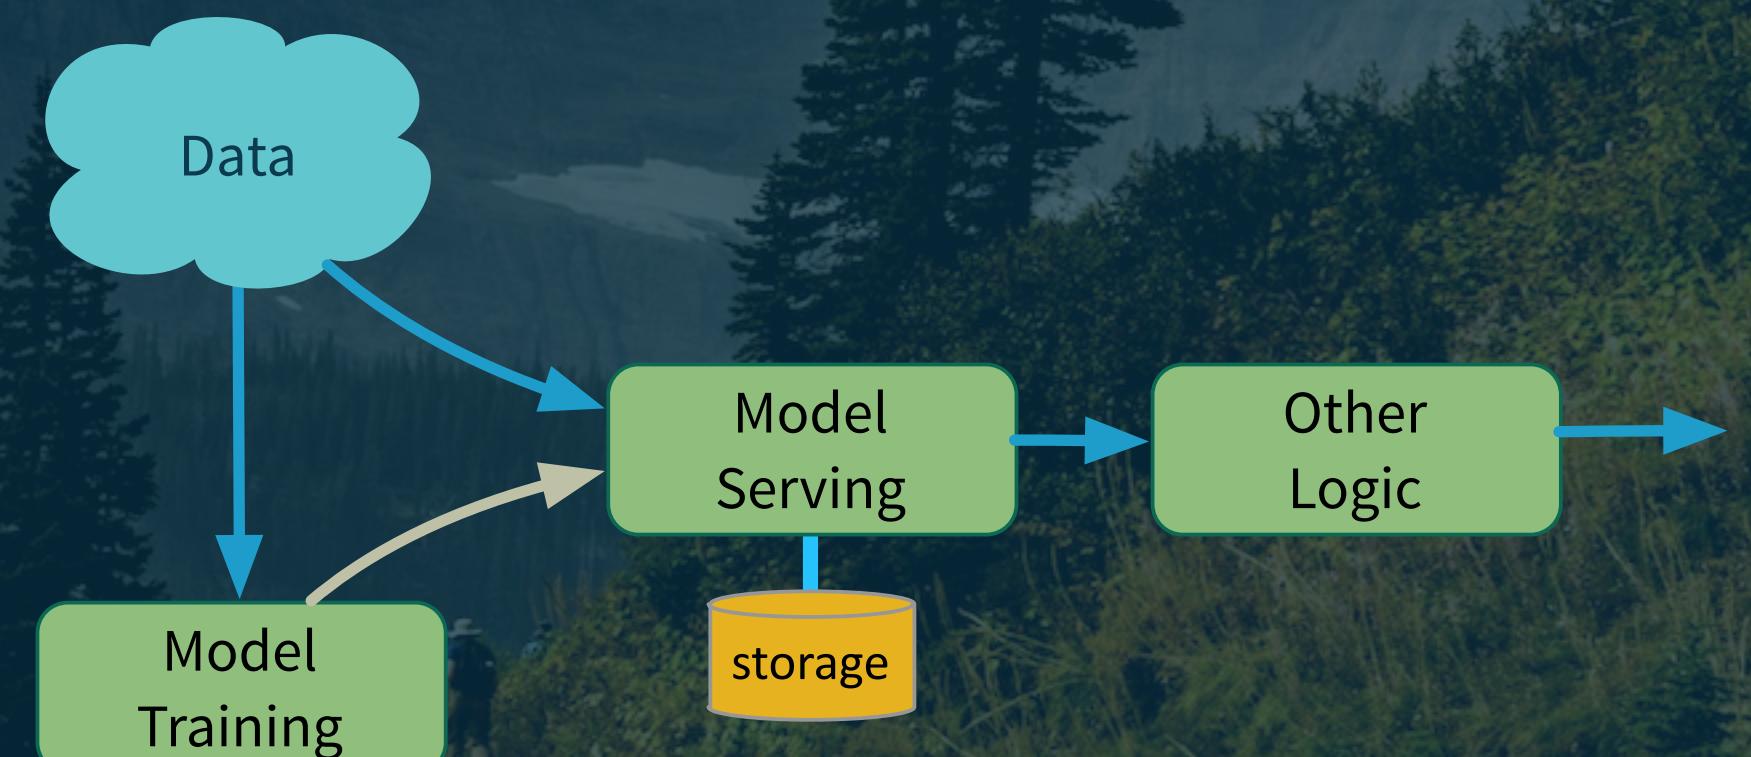
- Each datum has an identity
- Process each one uniquely
- Think sessions and state machines

A Spectrum of Microservices

Event-driven μ-services



“Record-centric” μ-services



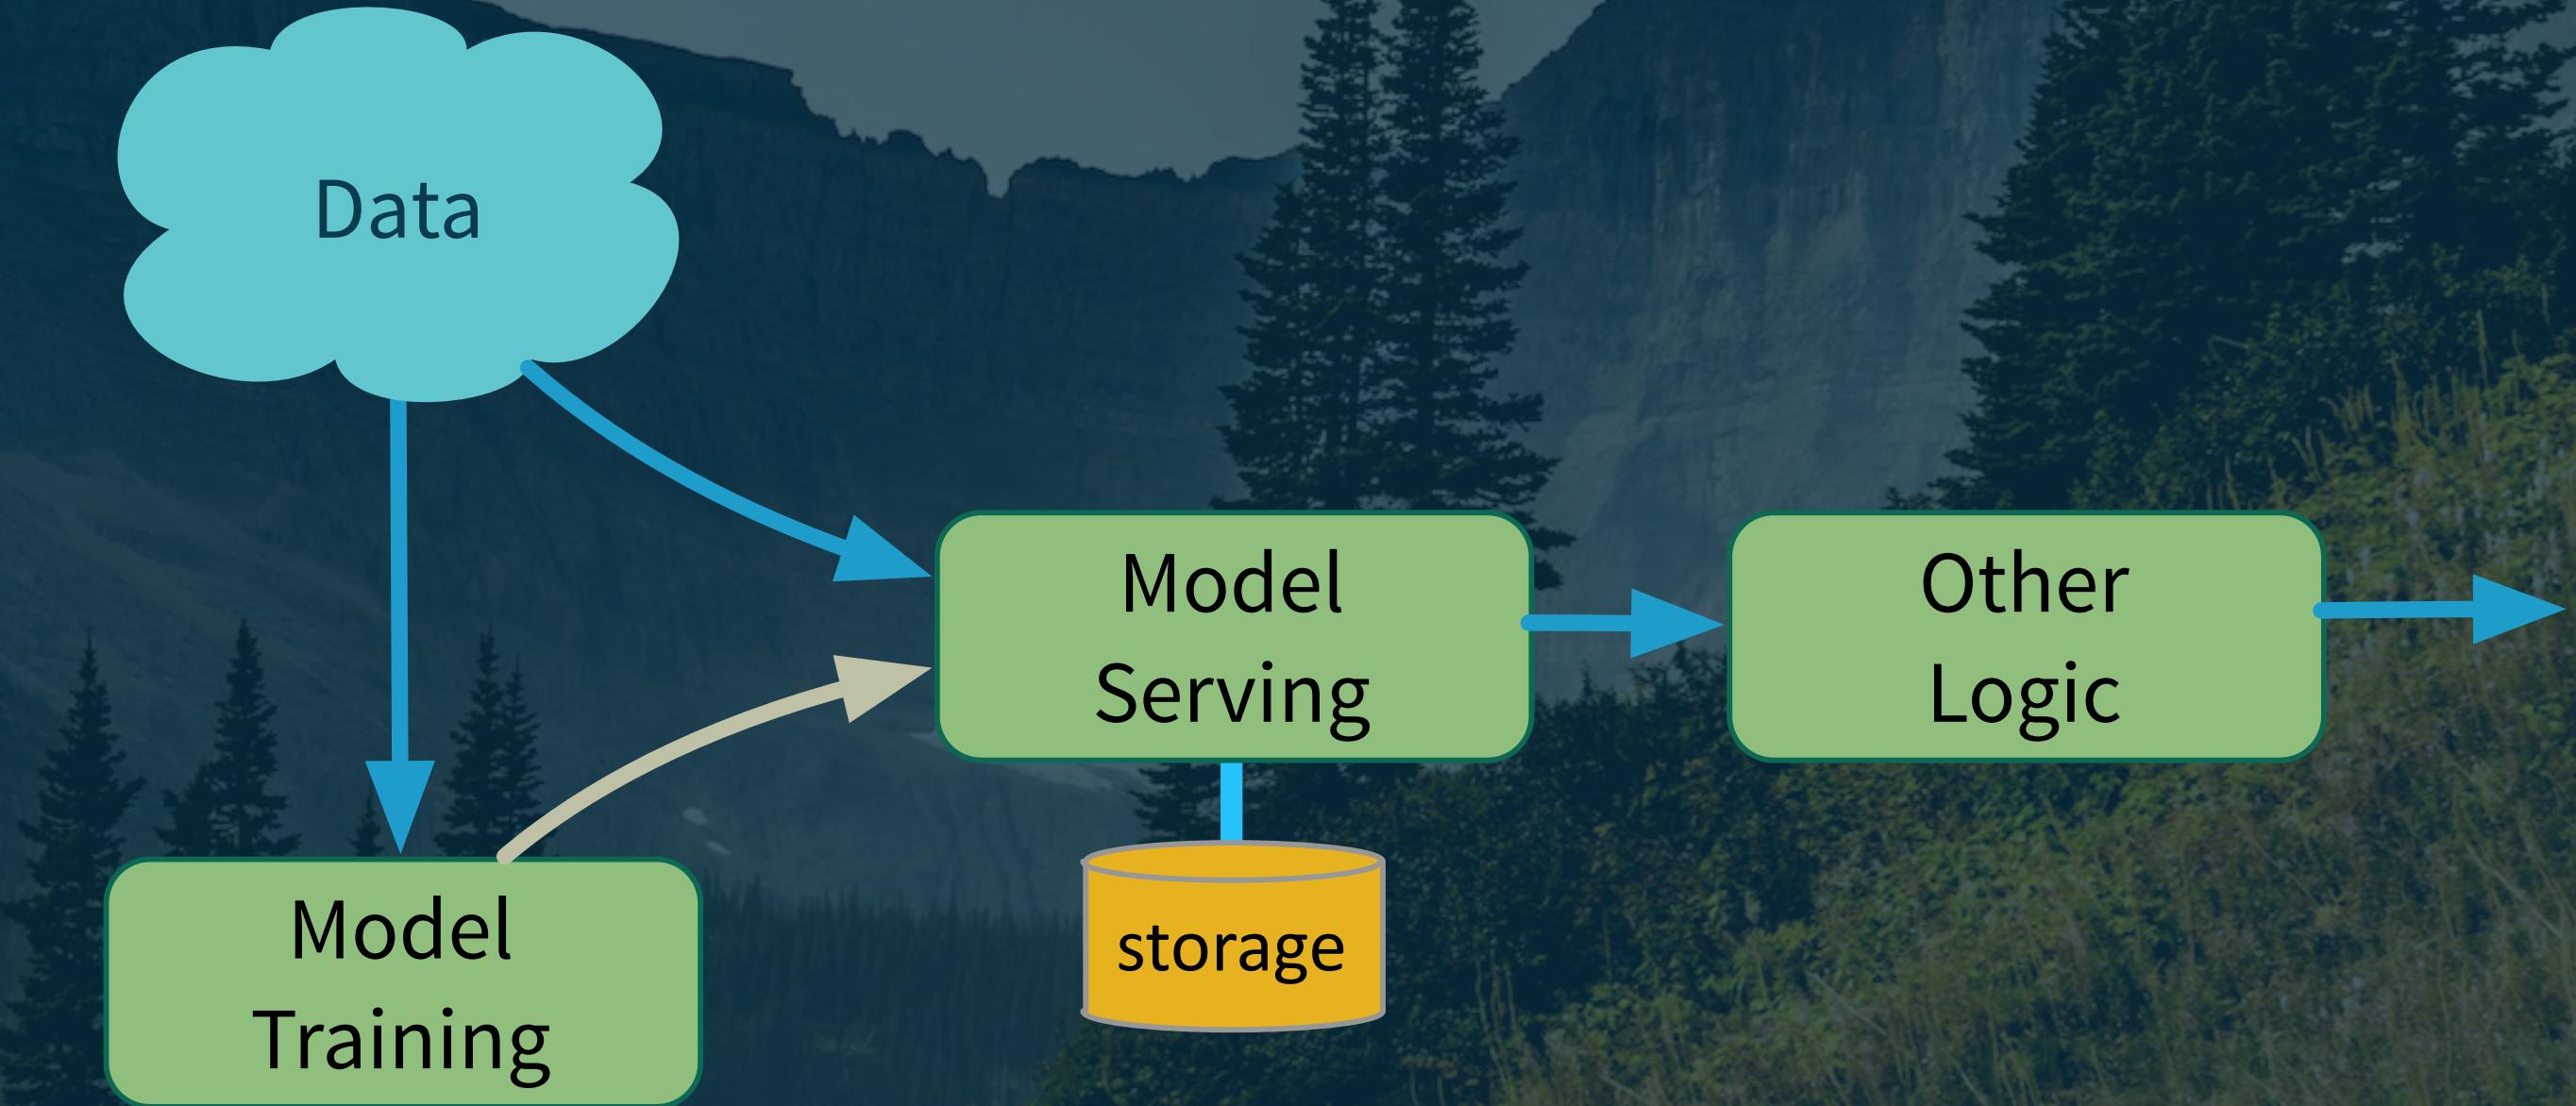
Events

← Data Spectrum →

Records

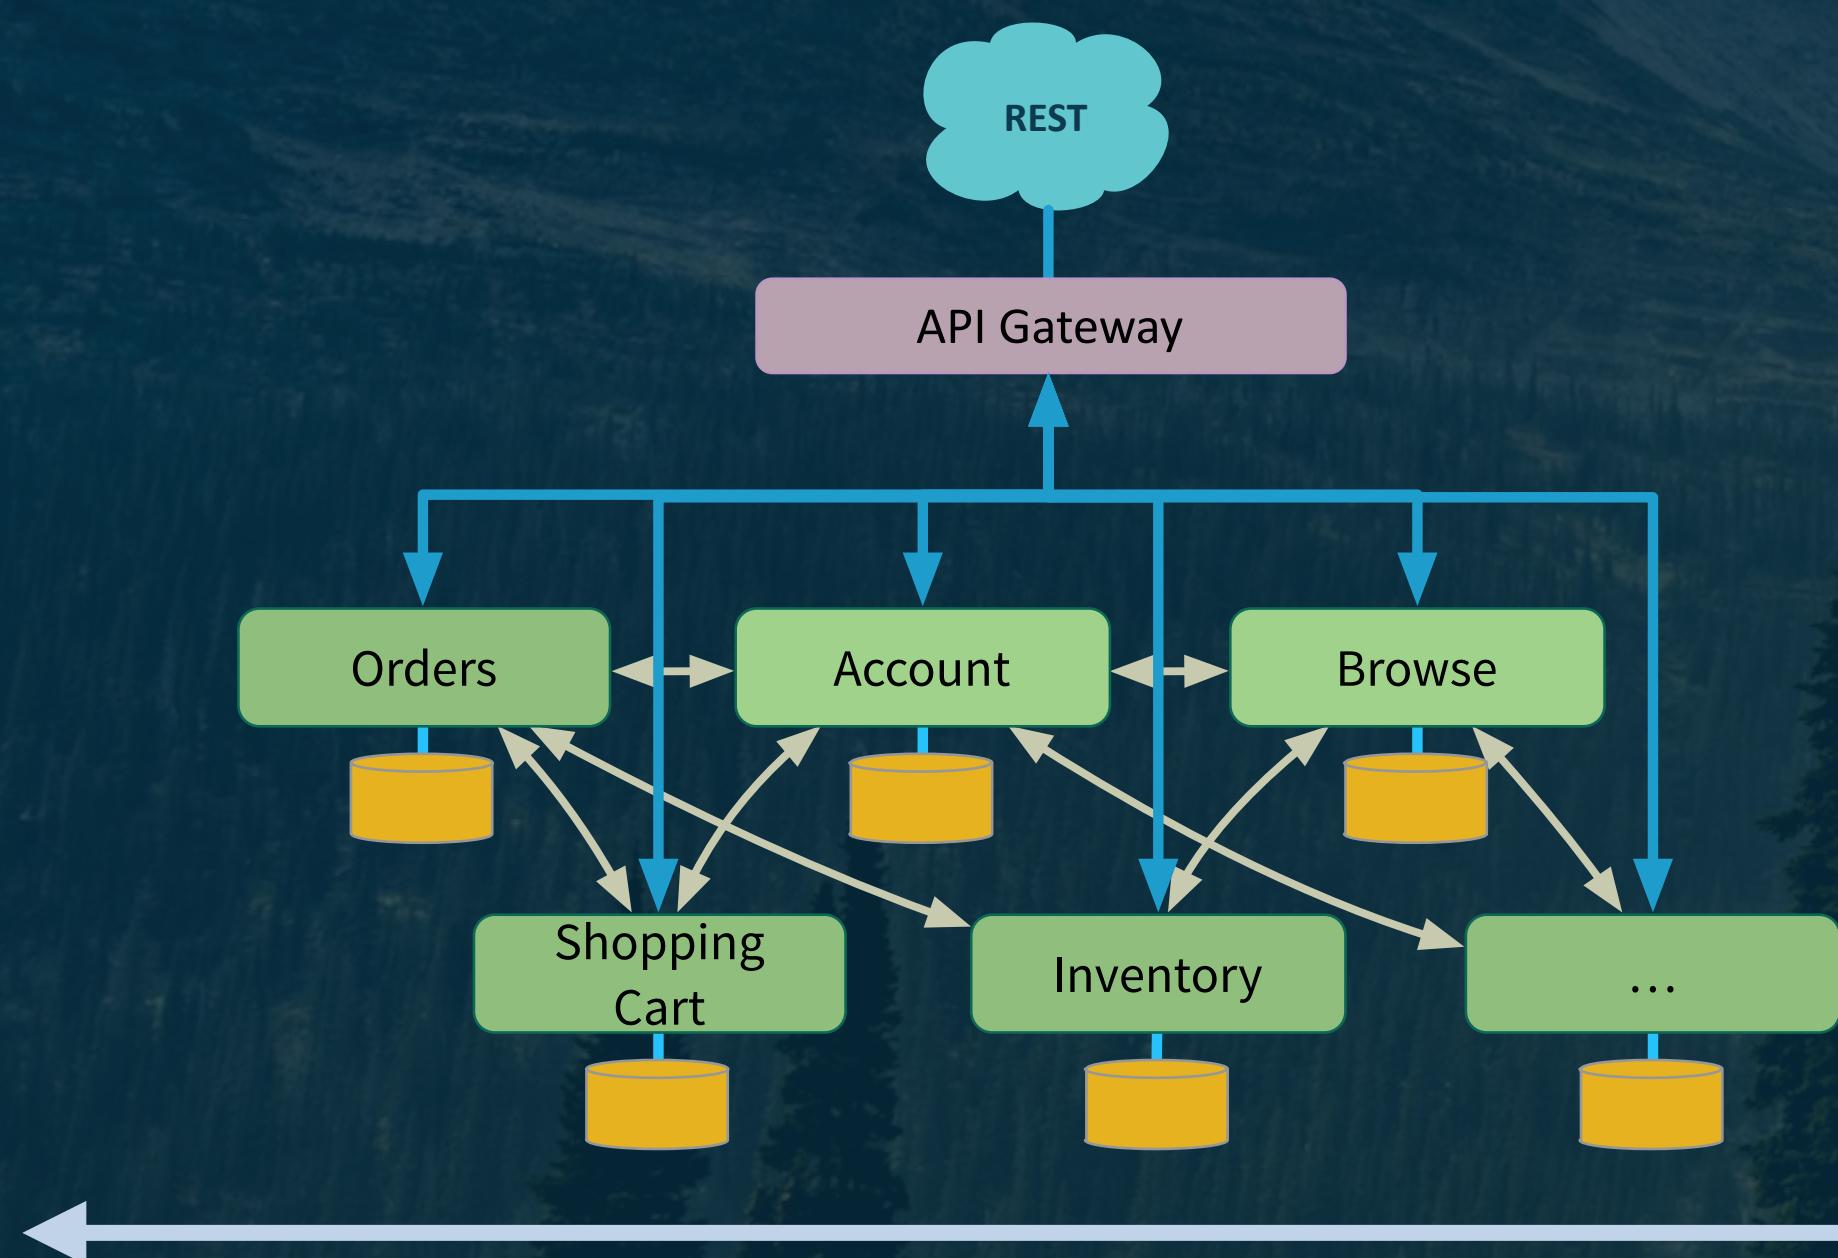
A Spectrum of Microservices

- “Anonymous” records
- Process *en masse*
- Think SQL queries for analytics

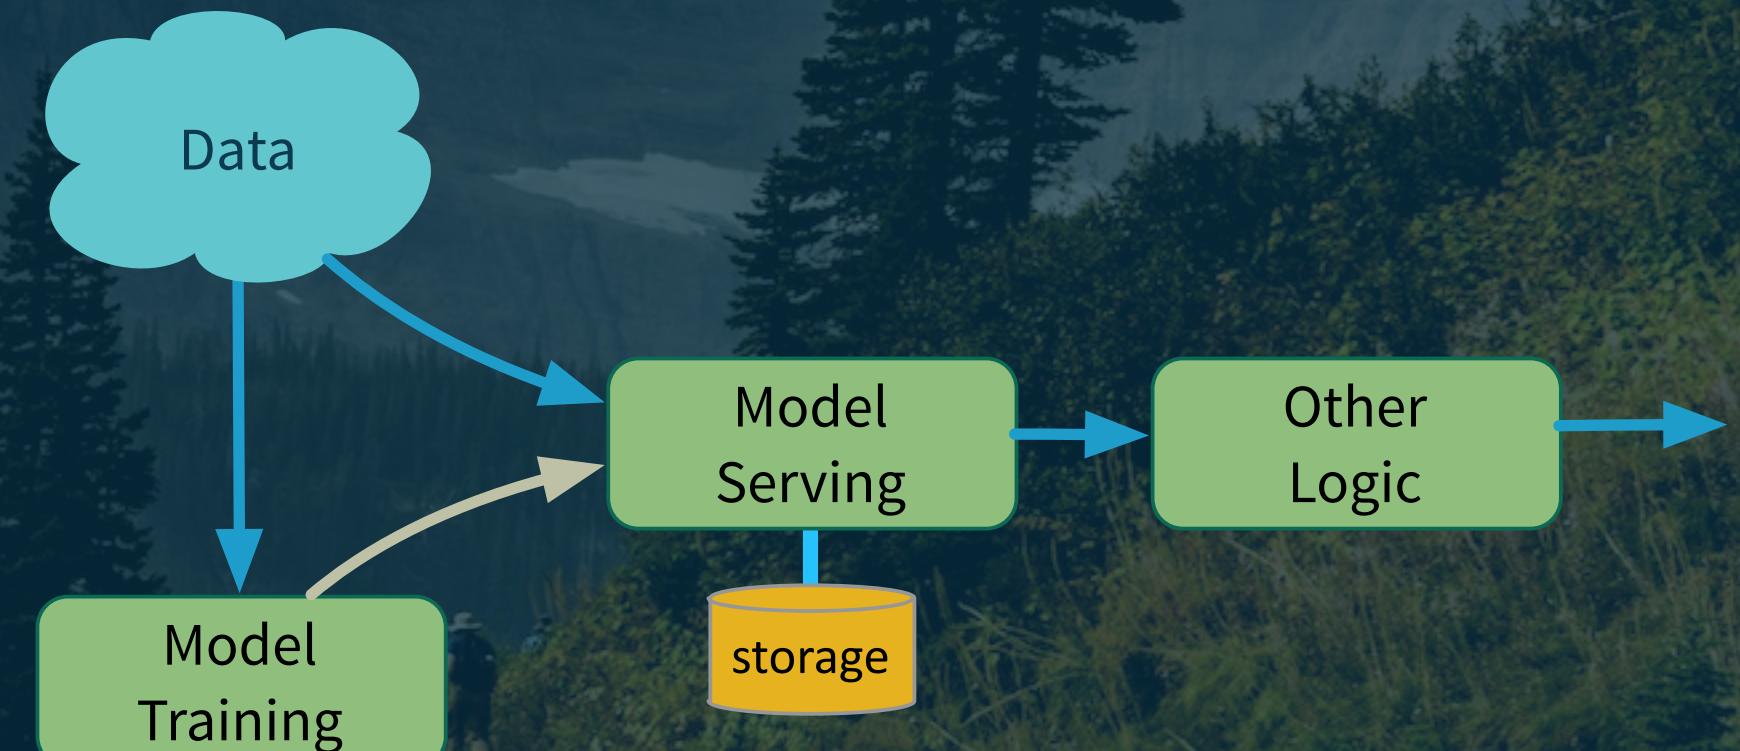


A Spectrum of Microservices

Event-driven μ-services



“Record-centric” μ-services



Events

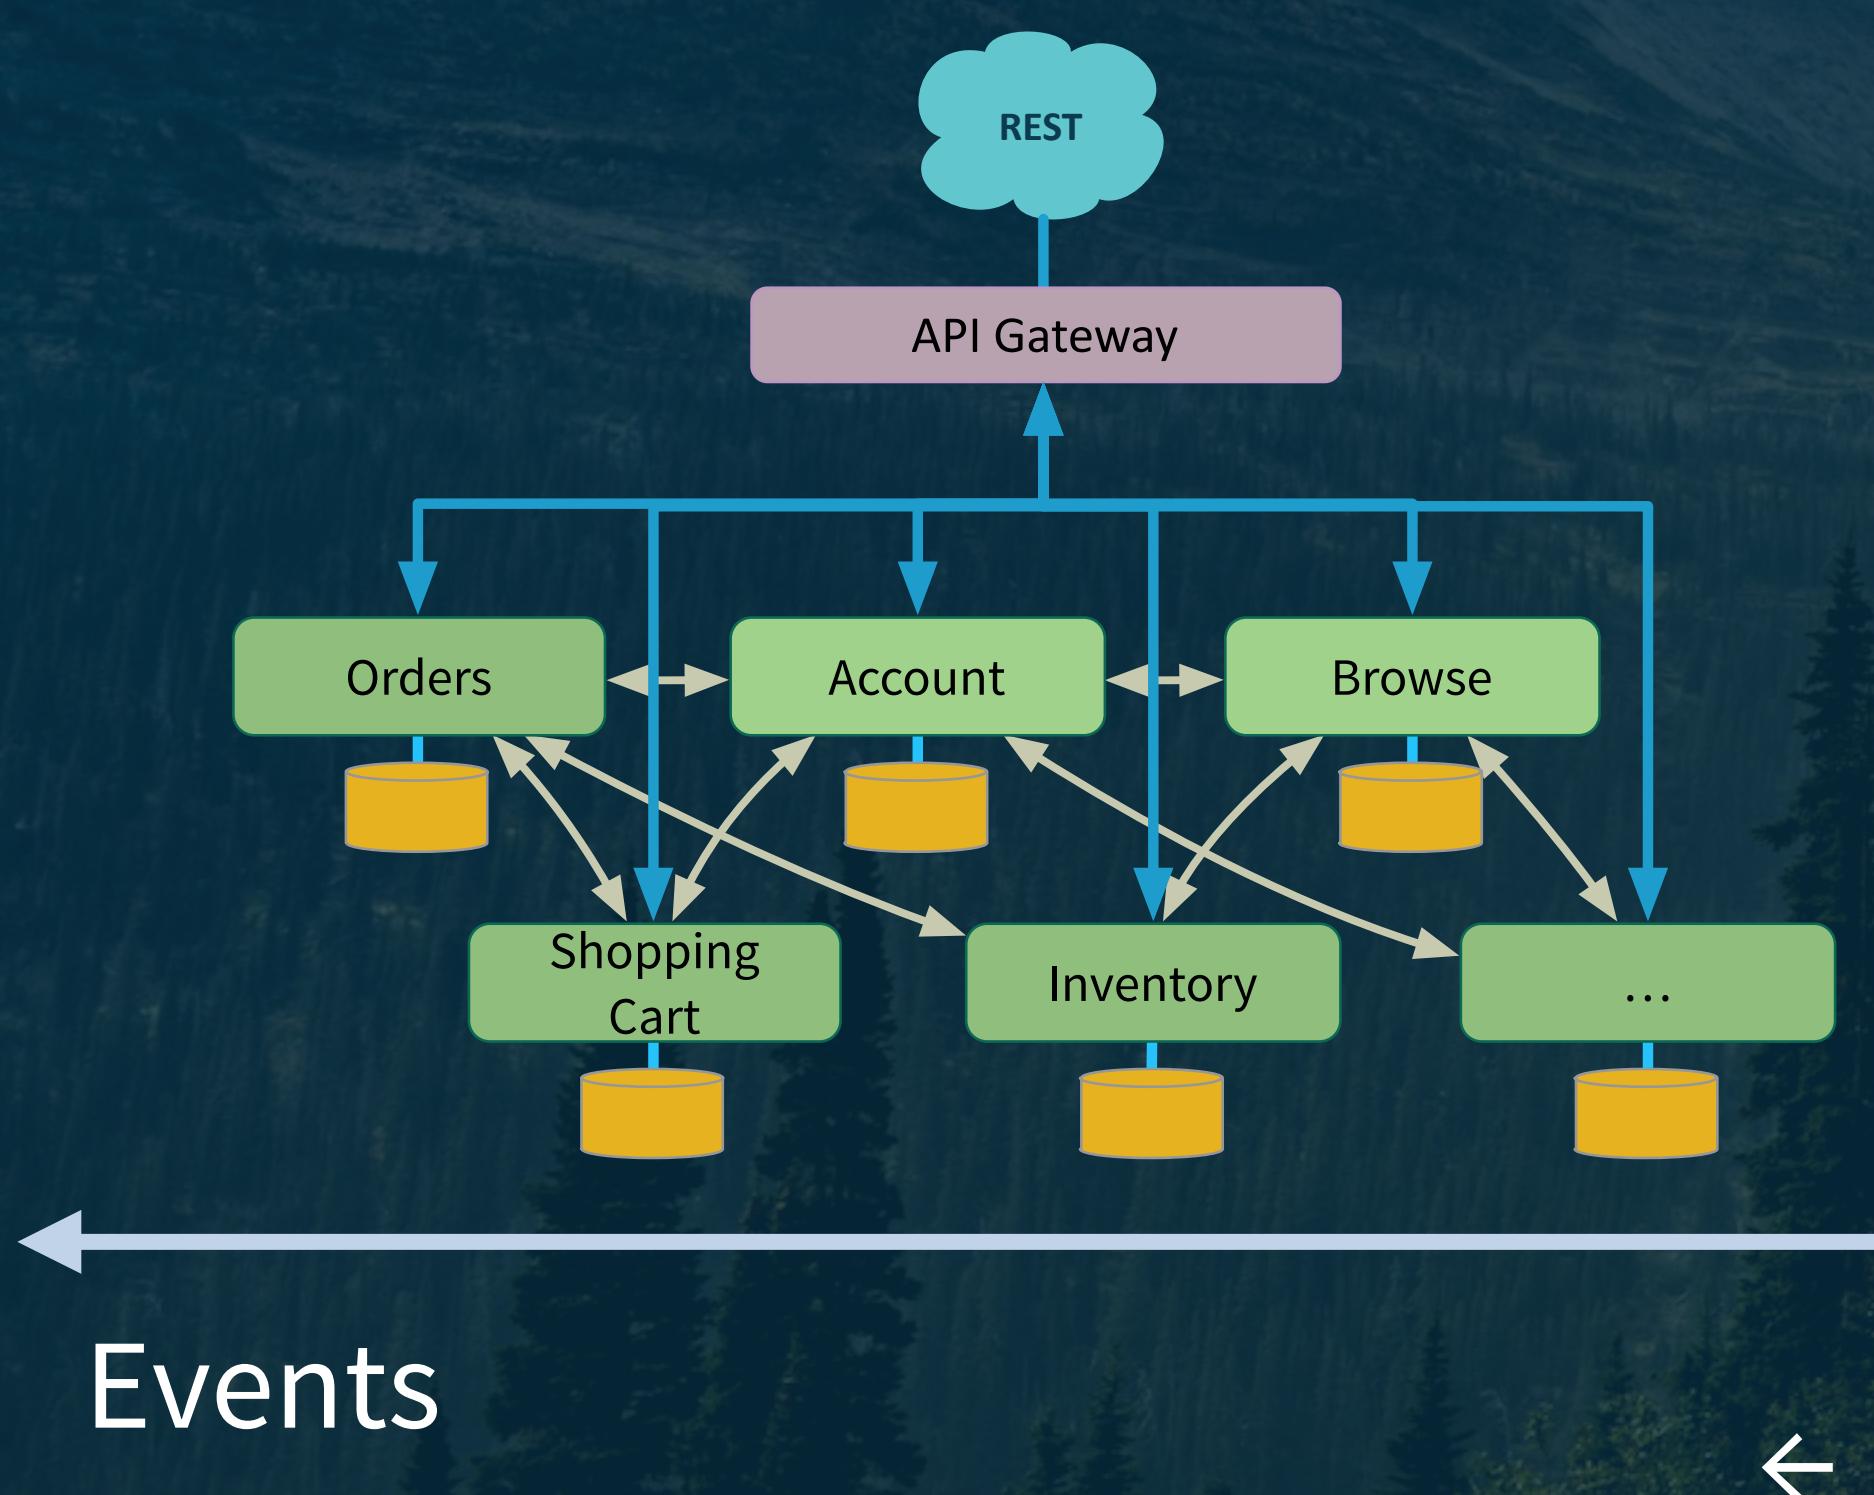
← Data Spectrum →

Records

A Spectrum of Microservices



Event-driven μ-services



Akka emerged from the left-hand side of the spectrum, the world of highly *Reactive* microservices.

Akka Streams pushes to the right, more data-centric.

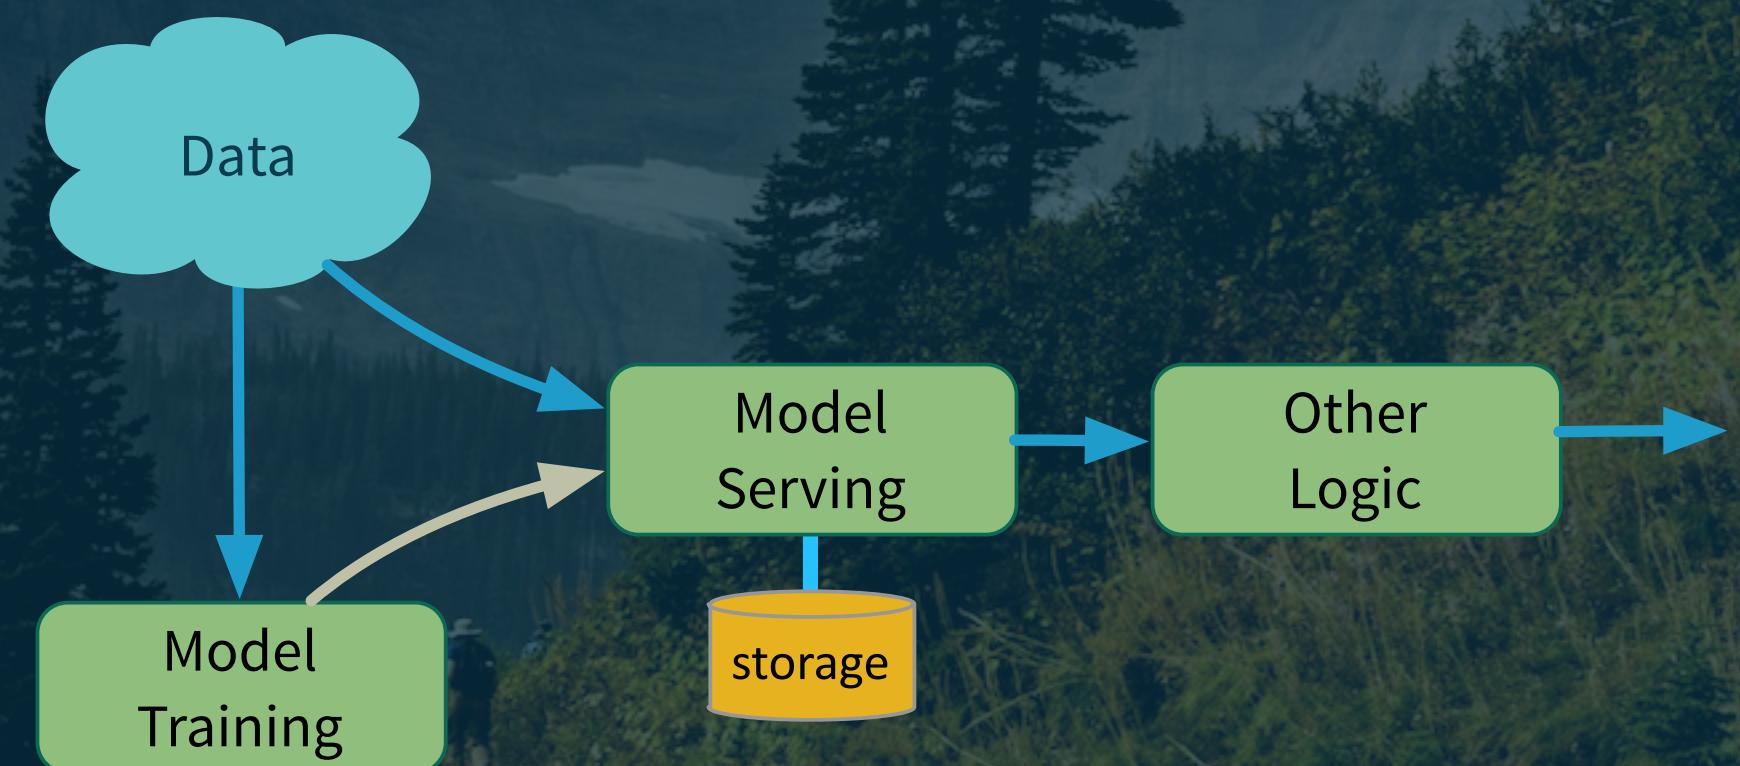
<https://www.reactivemanifesto.org/>

A Spectrum of Microservices



Emerged from the right-hand “Record-centric” μ -services side.

Kafka Streams pushes to the left, supporting many event-processing scenarios.



Events

← Data Spectrum →

Records



kafka

Kafka Streams

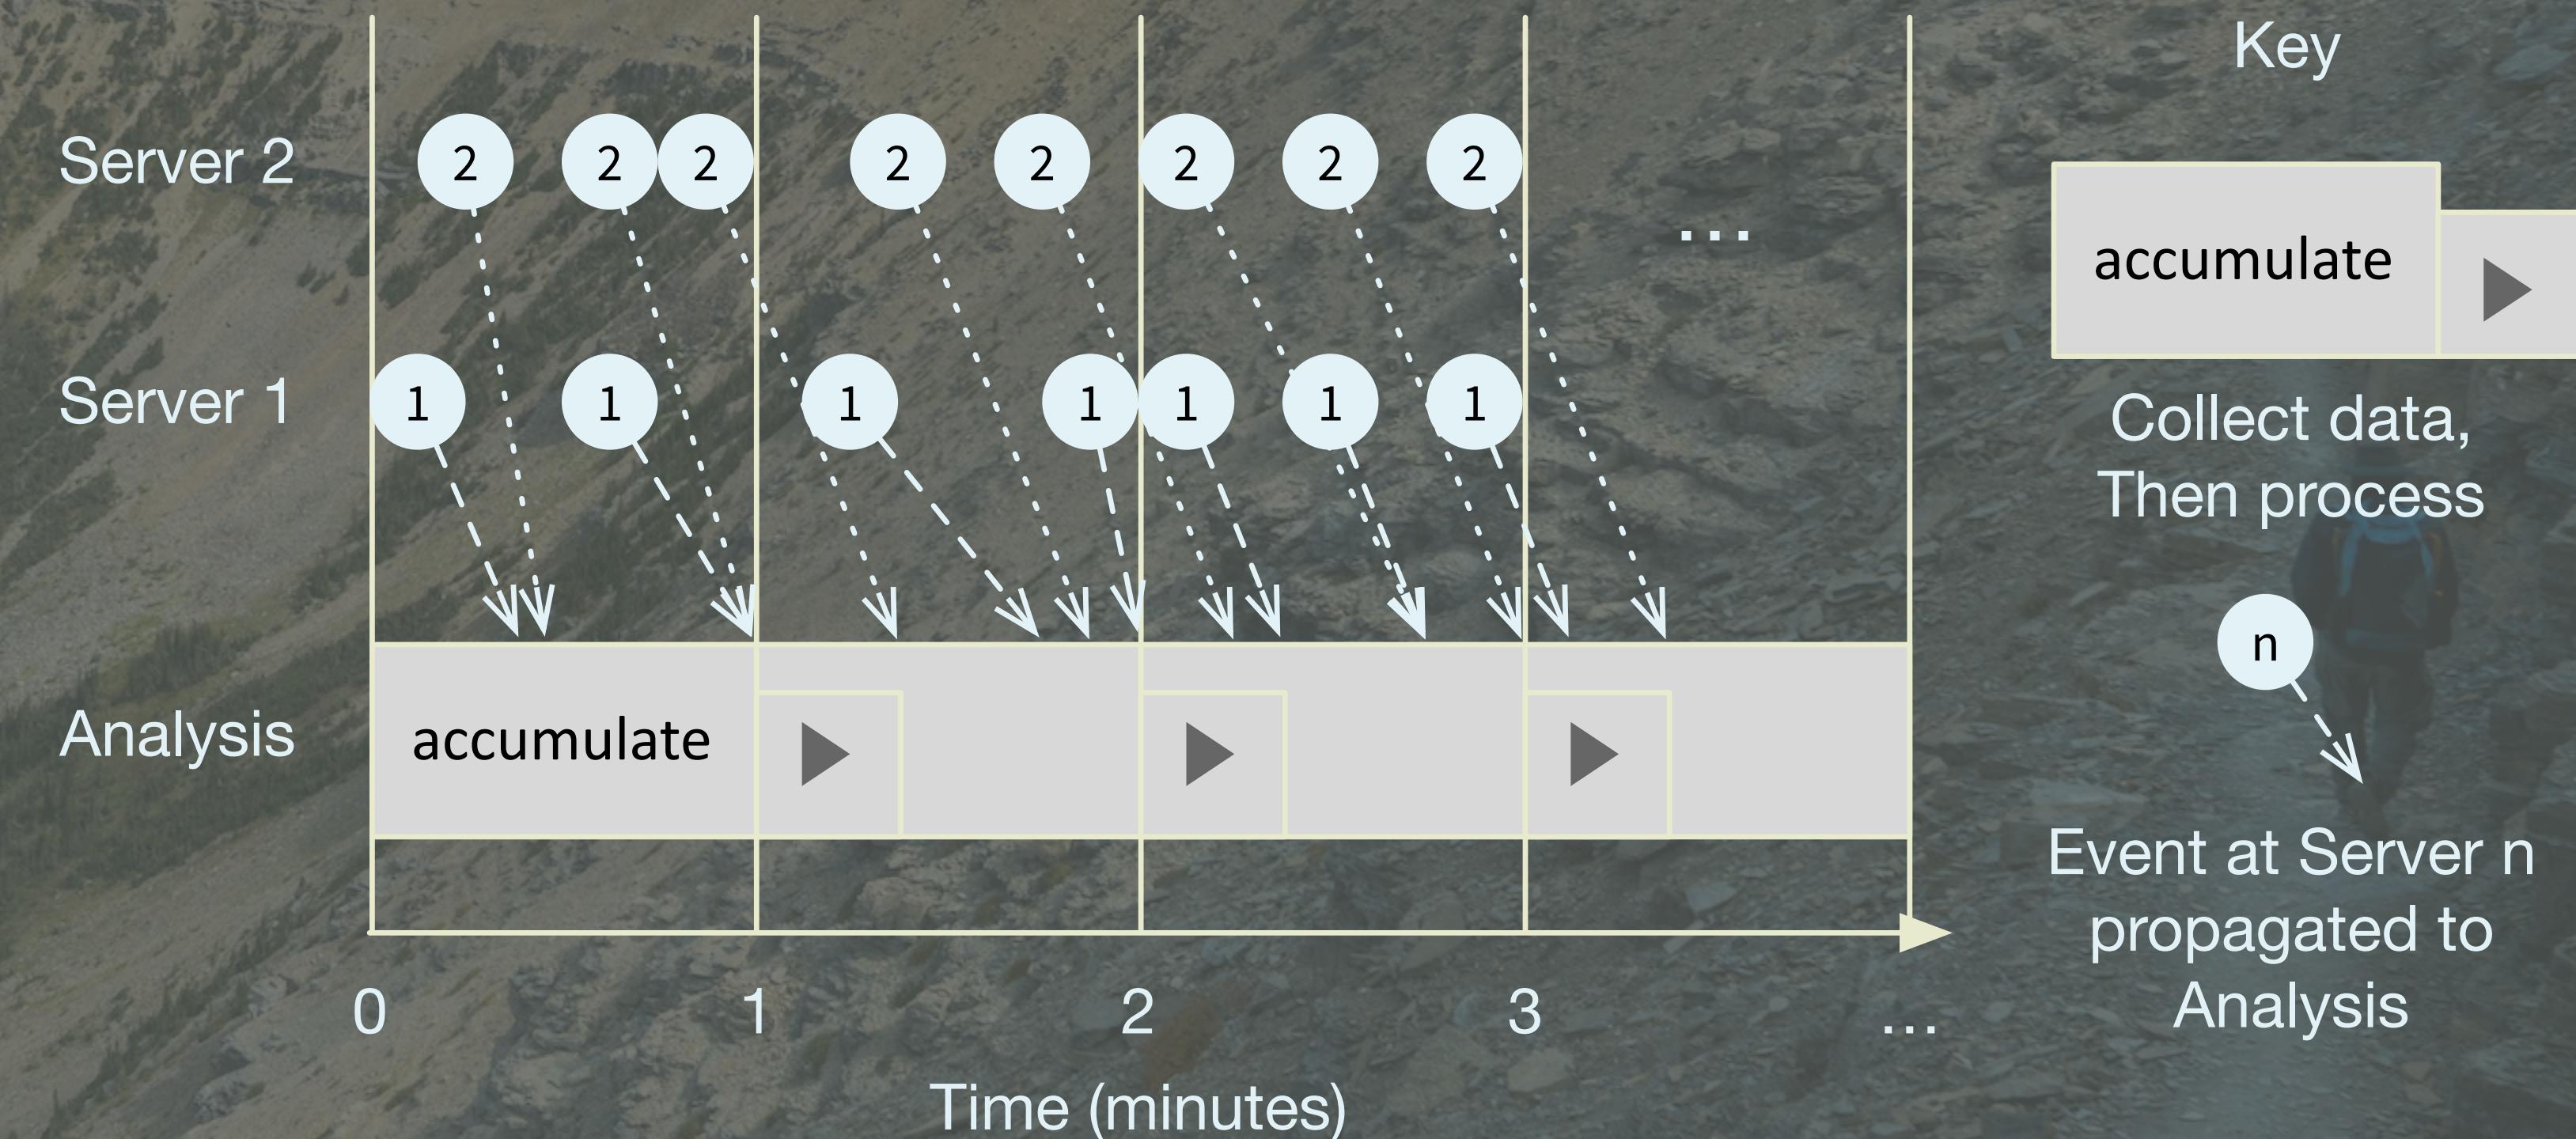




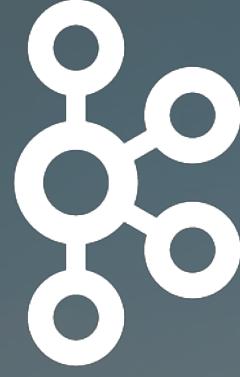
kafka

Kafka Streams

- Important stream-processing semantics, e.g.,
- Windowing support (e.g., group by within a window)



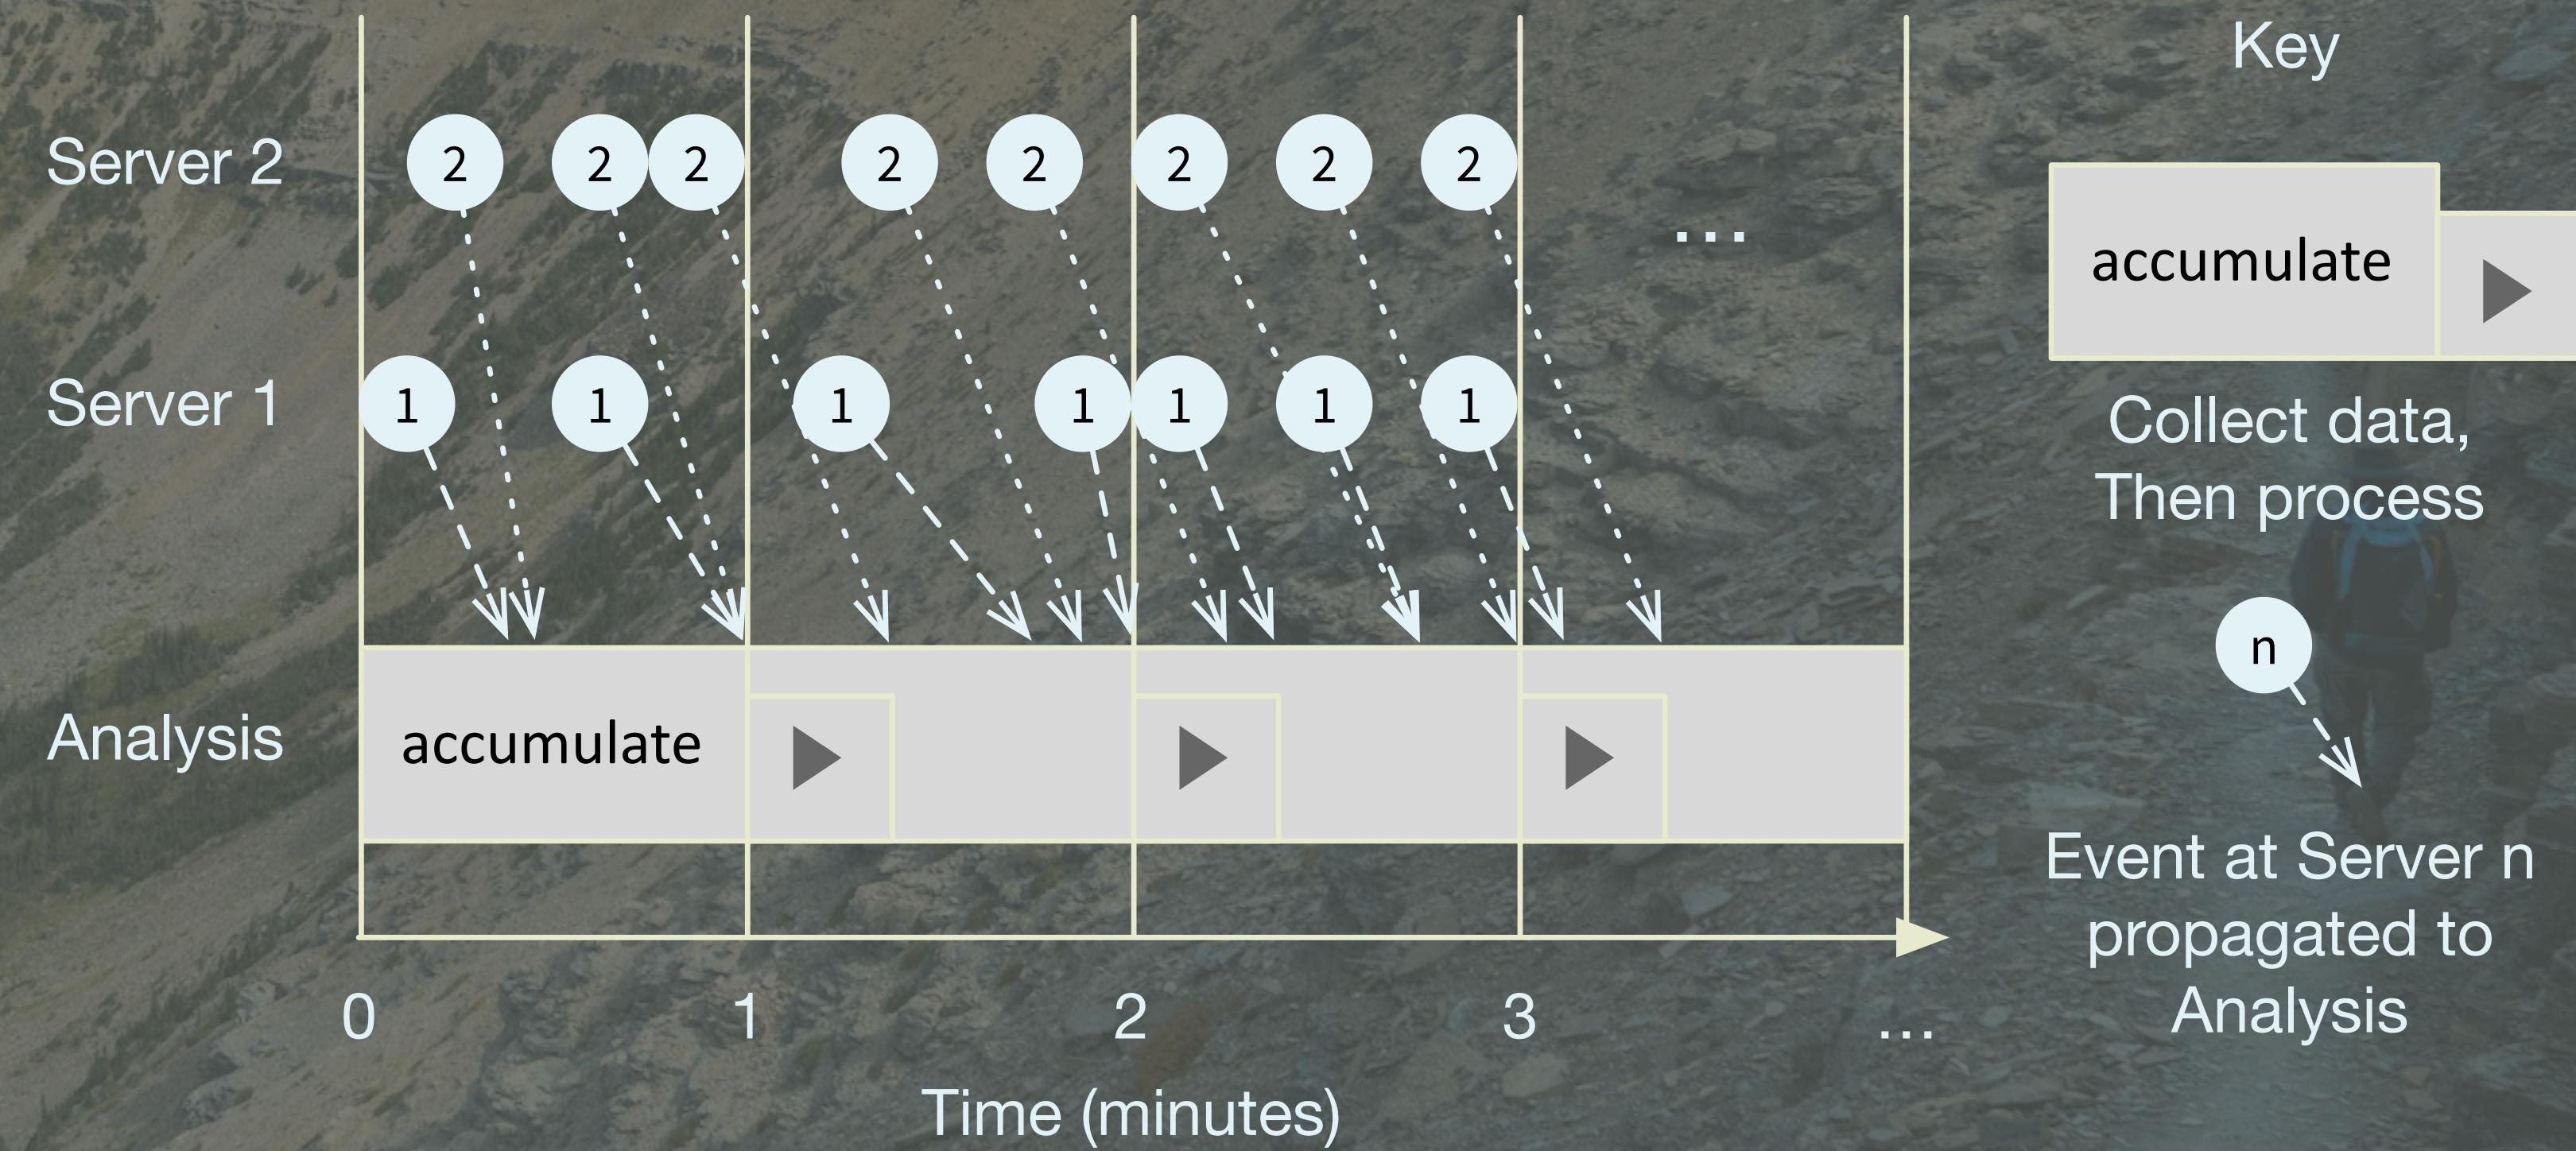
See my O'Reilly
report for details.



kafka

Kafka Streams

- Important stream-processing semantics, e.g.,
- Distinguish between *event time* and *processing time*





kafka

Kafka Streams

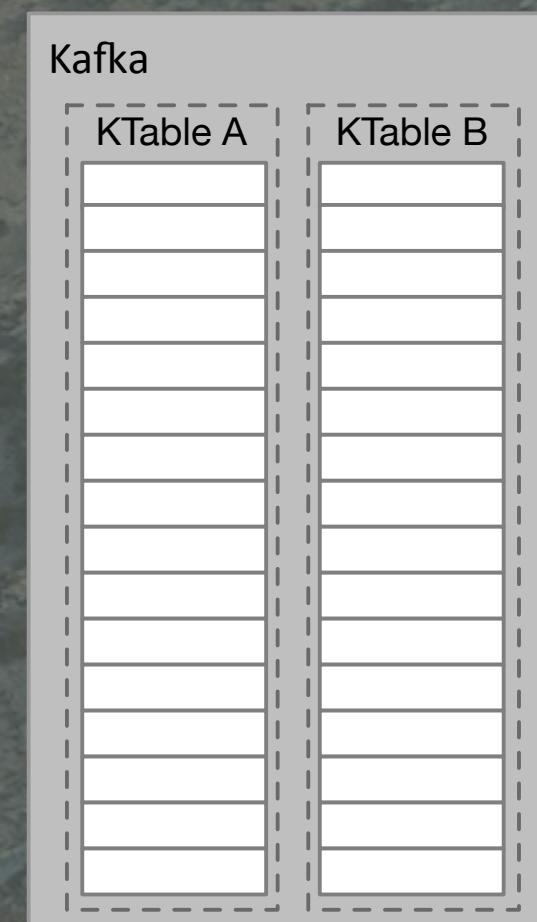
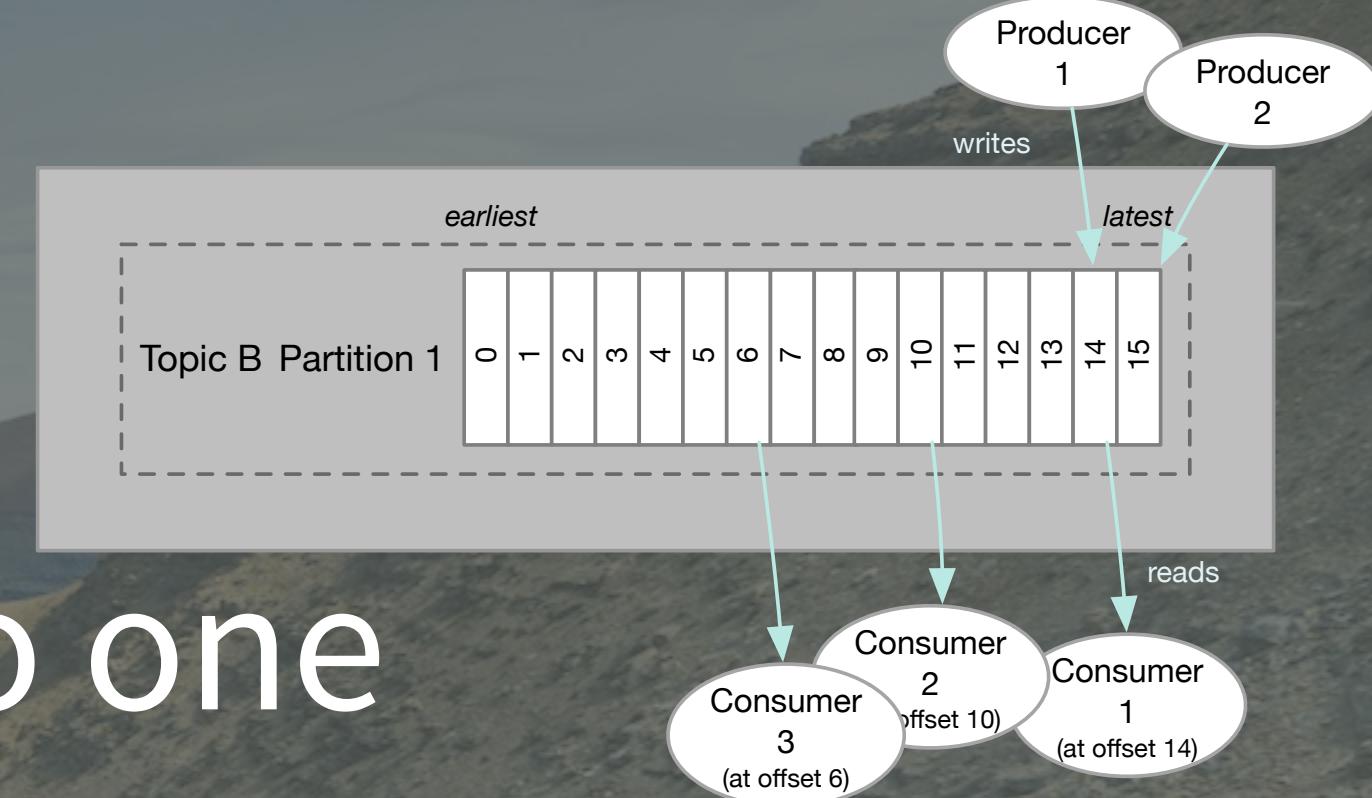
- *Exactly once*
- vs.
 - *At most once* - “fire and forget”
 - *At least once* - keep trying until acknowledged
(but you have to handle de-duplication)
- Note: it’s really *effectively once*



kafka

Kafka Streams

- KStream
 - Per record transformations, one to one mapping
- KTable
 - Last value per key
 - Useful for holding *state* values





Kafka Streams

kafka

- Read from and write to Kafka topics, memory
 - Kafka Connect used for other sources and sinks
- Load balance and scale using topic partitioning

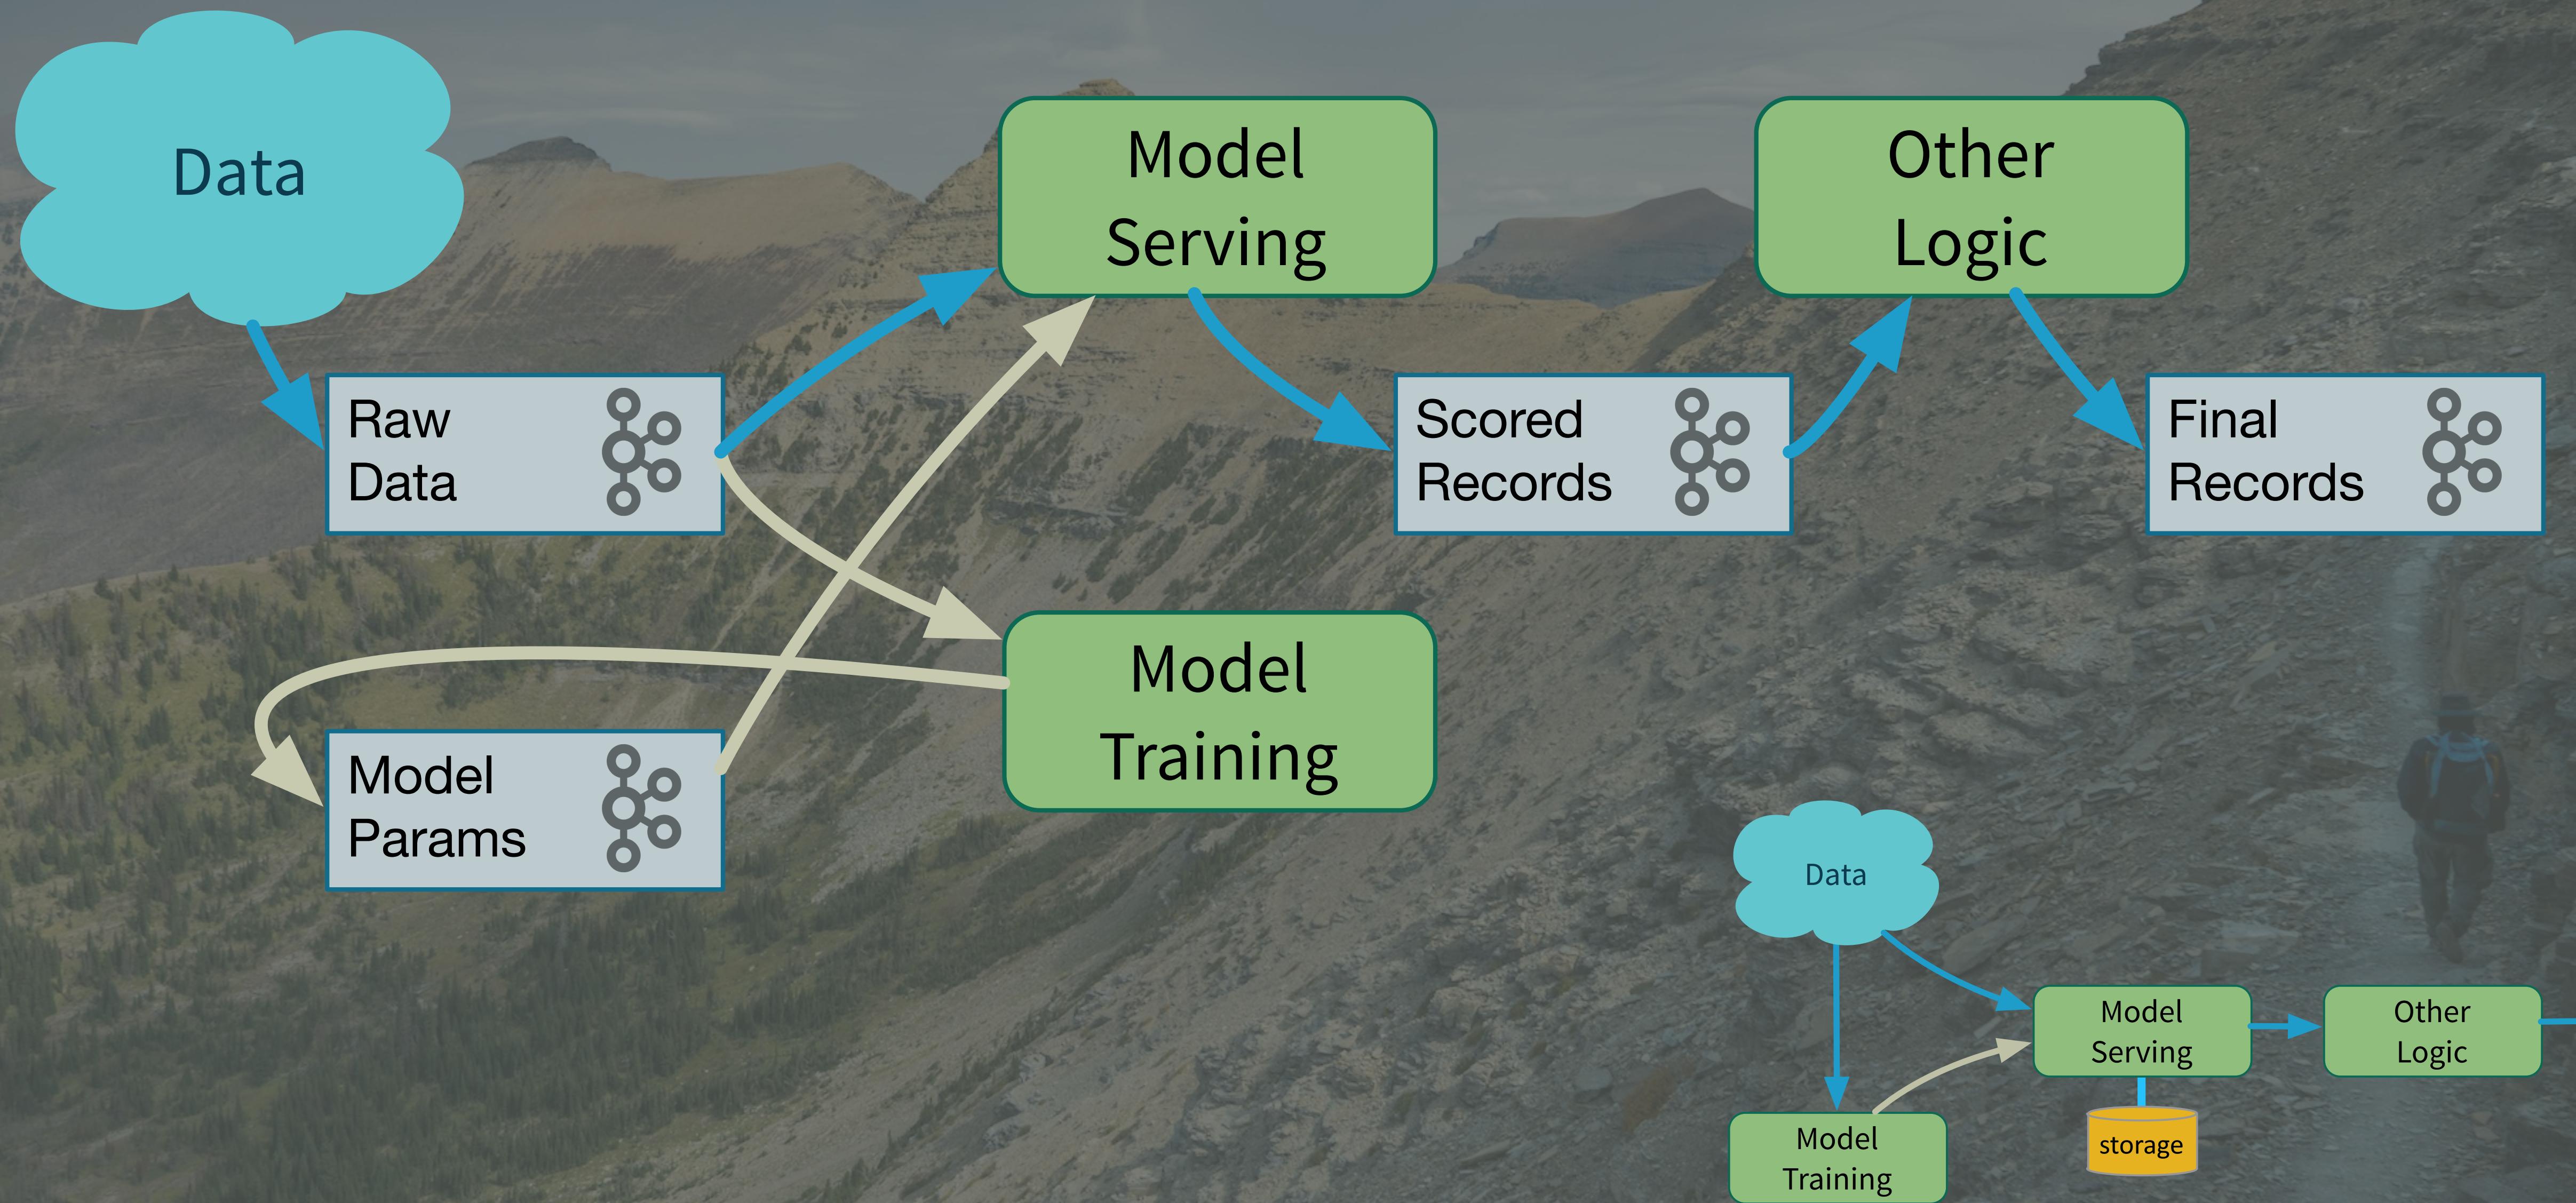


Kafka Streams

kafka

- Java API
- Scala API: written by Lightbend
- SQL!!

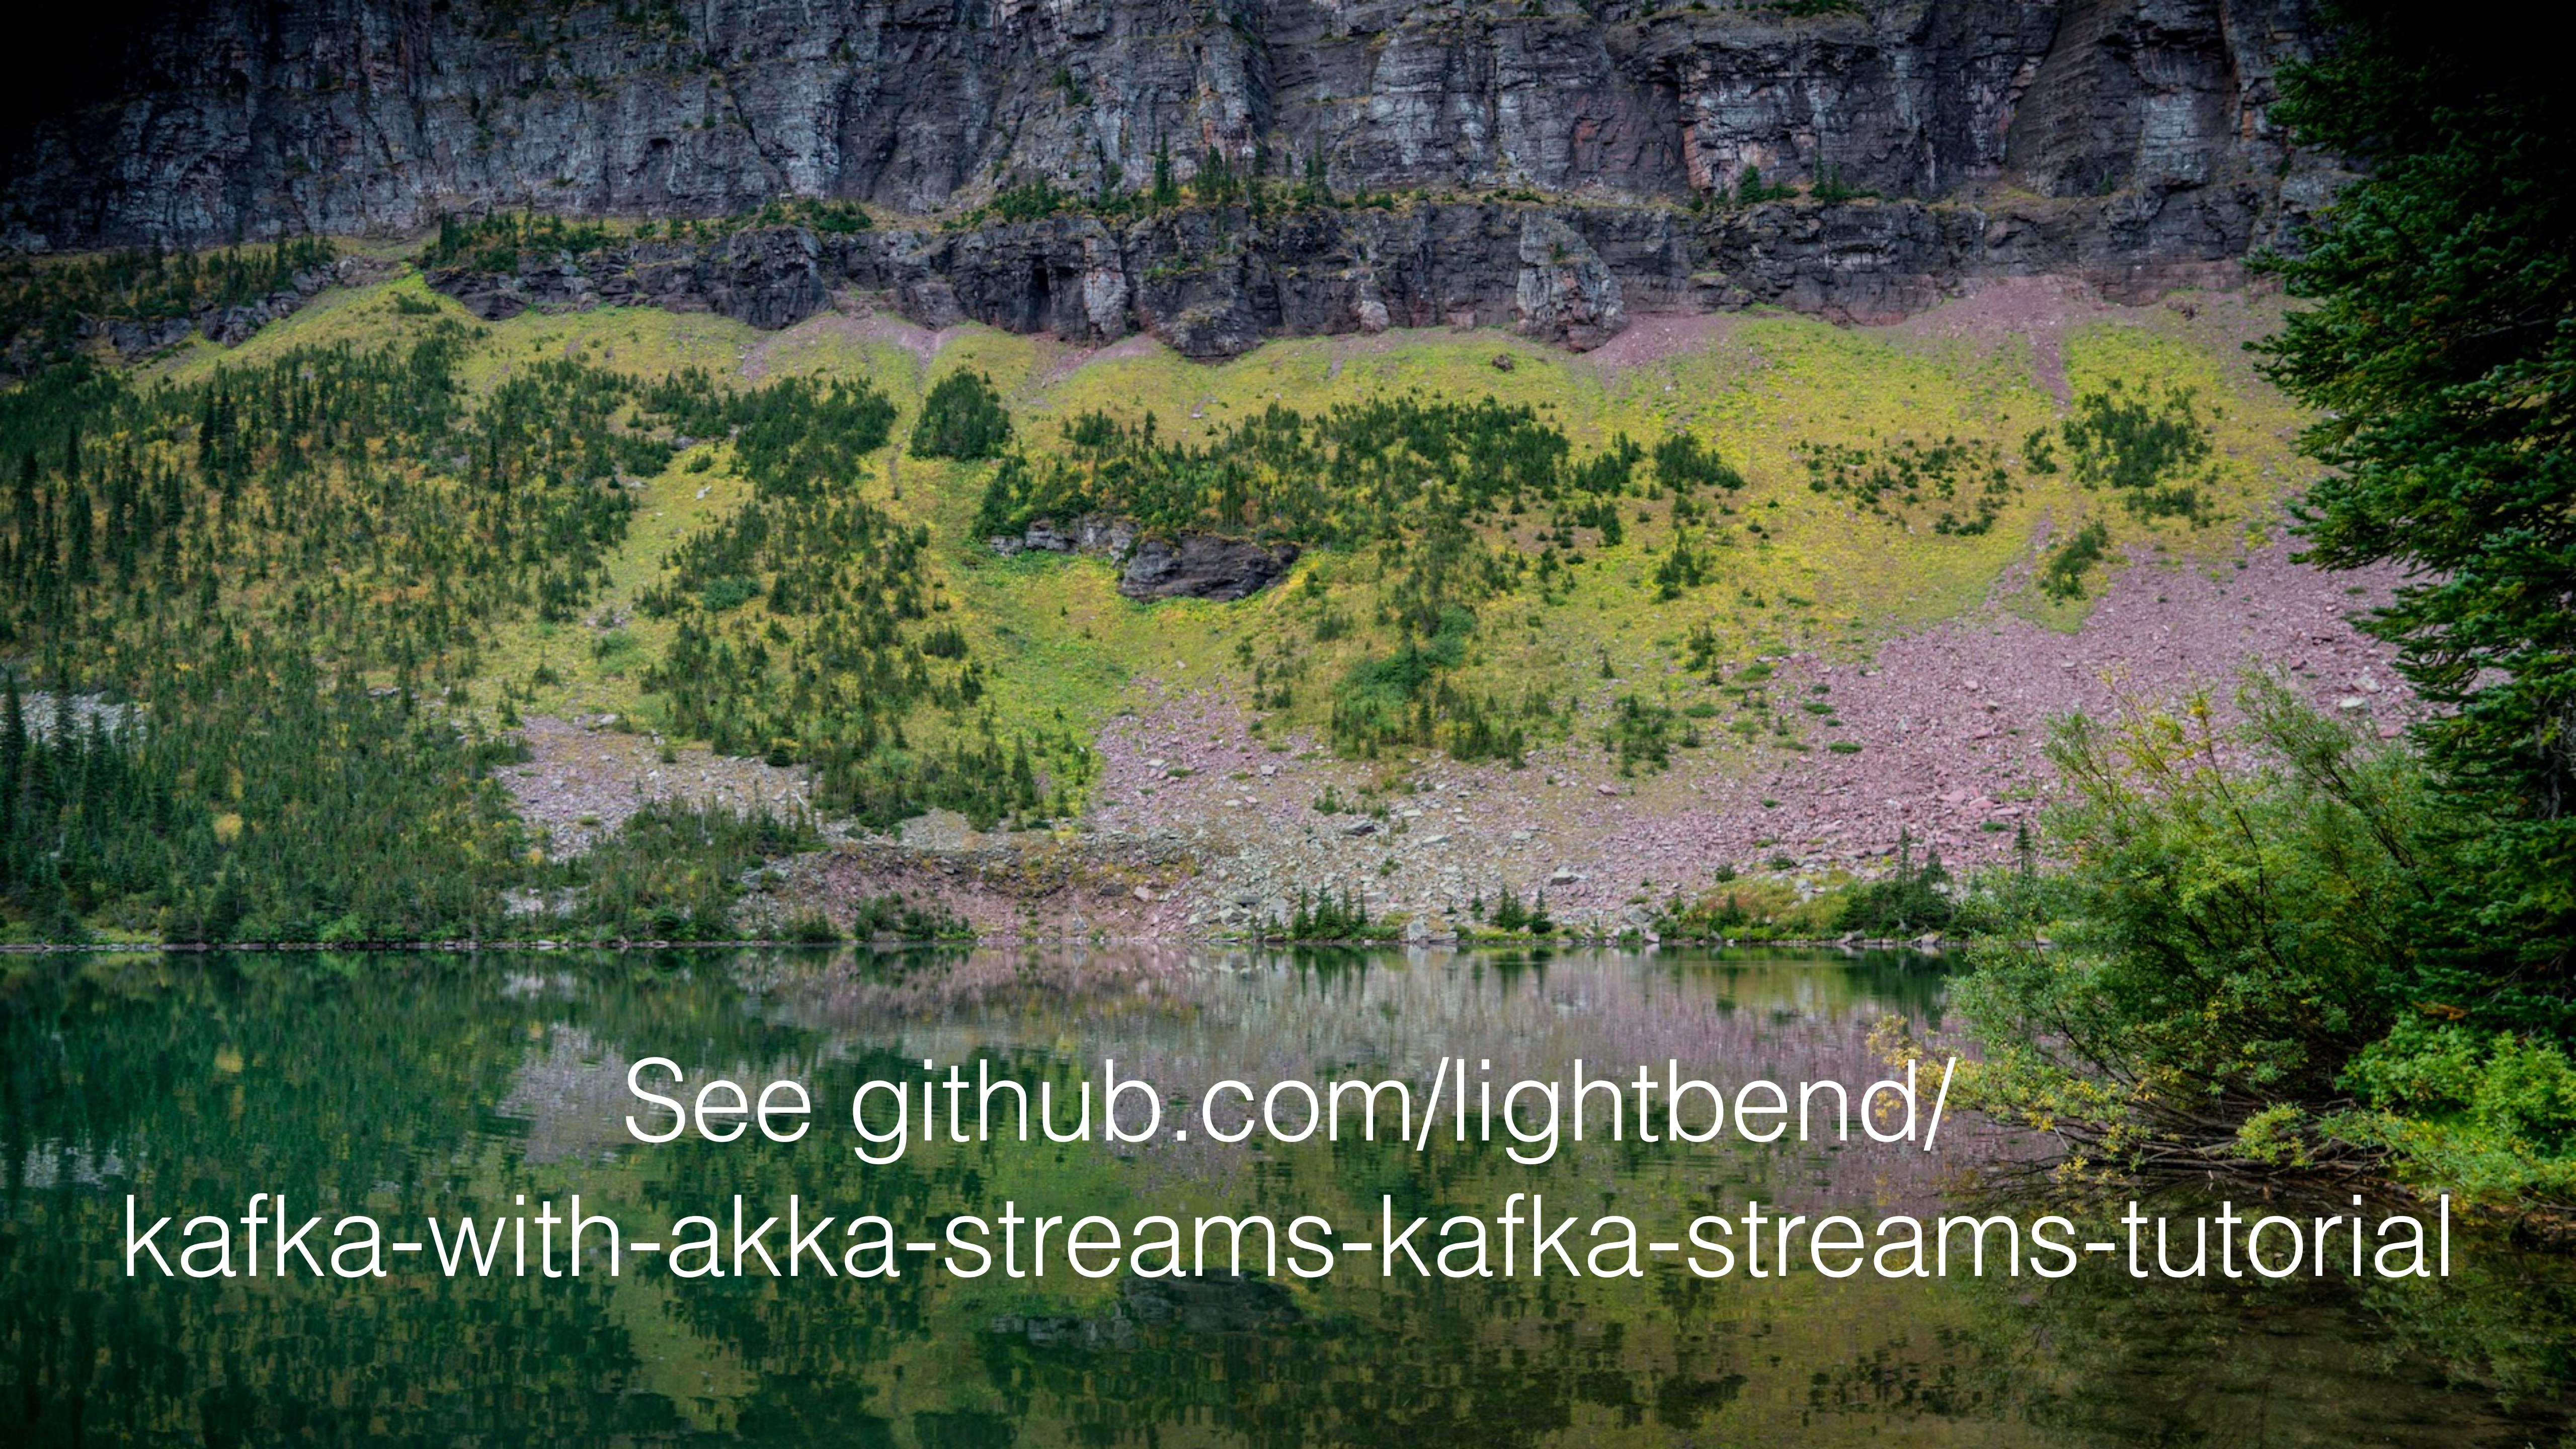
Kafka Streams Example



Kafka Streams Example

We'll use the new Scala Kafka Streams API:

- Developed by Debasish Ghosh, Boris Lublinsky, Sean Glover, et al. from Lightbend
- See also the following, if you know about *queryable state*:
 - <https://github.com/lightbend/kafka-streams-query>

A scenic landscape featuring a calm lake in the foreground, its surface reflecting the surrounding green hills and a large, dark, layered rock cliff above. The hills are covered with patches of green vegetation and some yellowish-green areas. A dense forest of evergreen trees is visible on the right side.

See [github.com/lightbend/
kafka-with-akka-streams-kafka-streams-tutorial](https://github.com/lightbend/kafka-with-akka-streams-kafka-streams-tutorial)

```

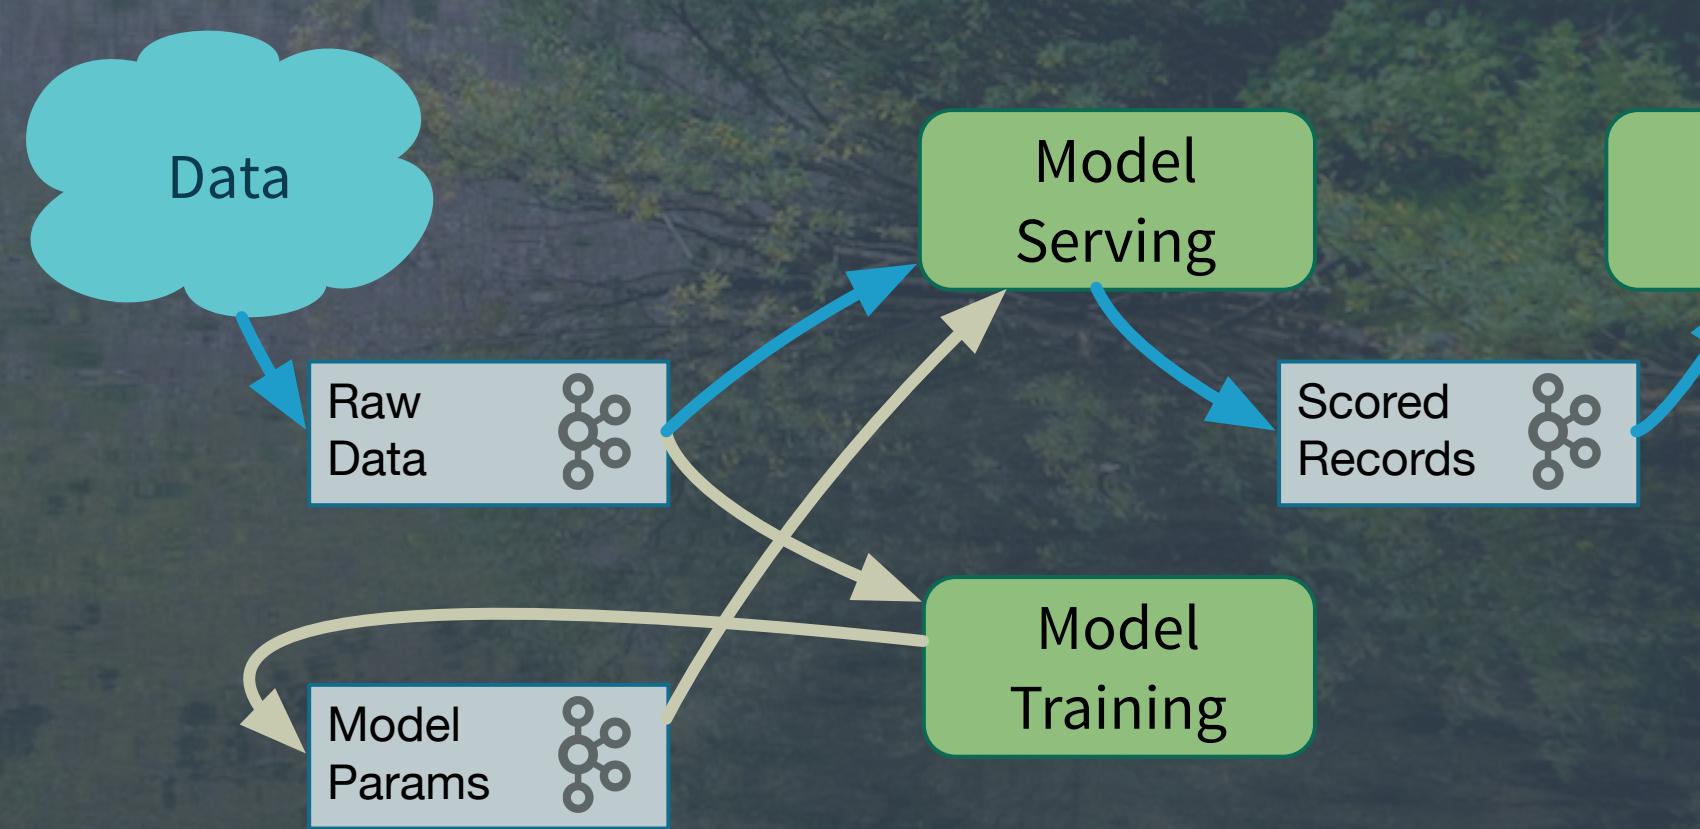
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```

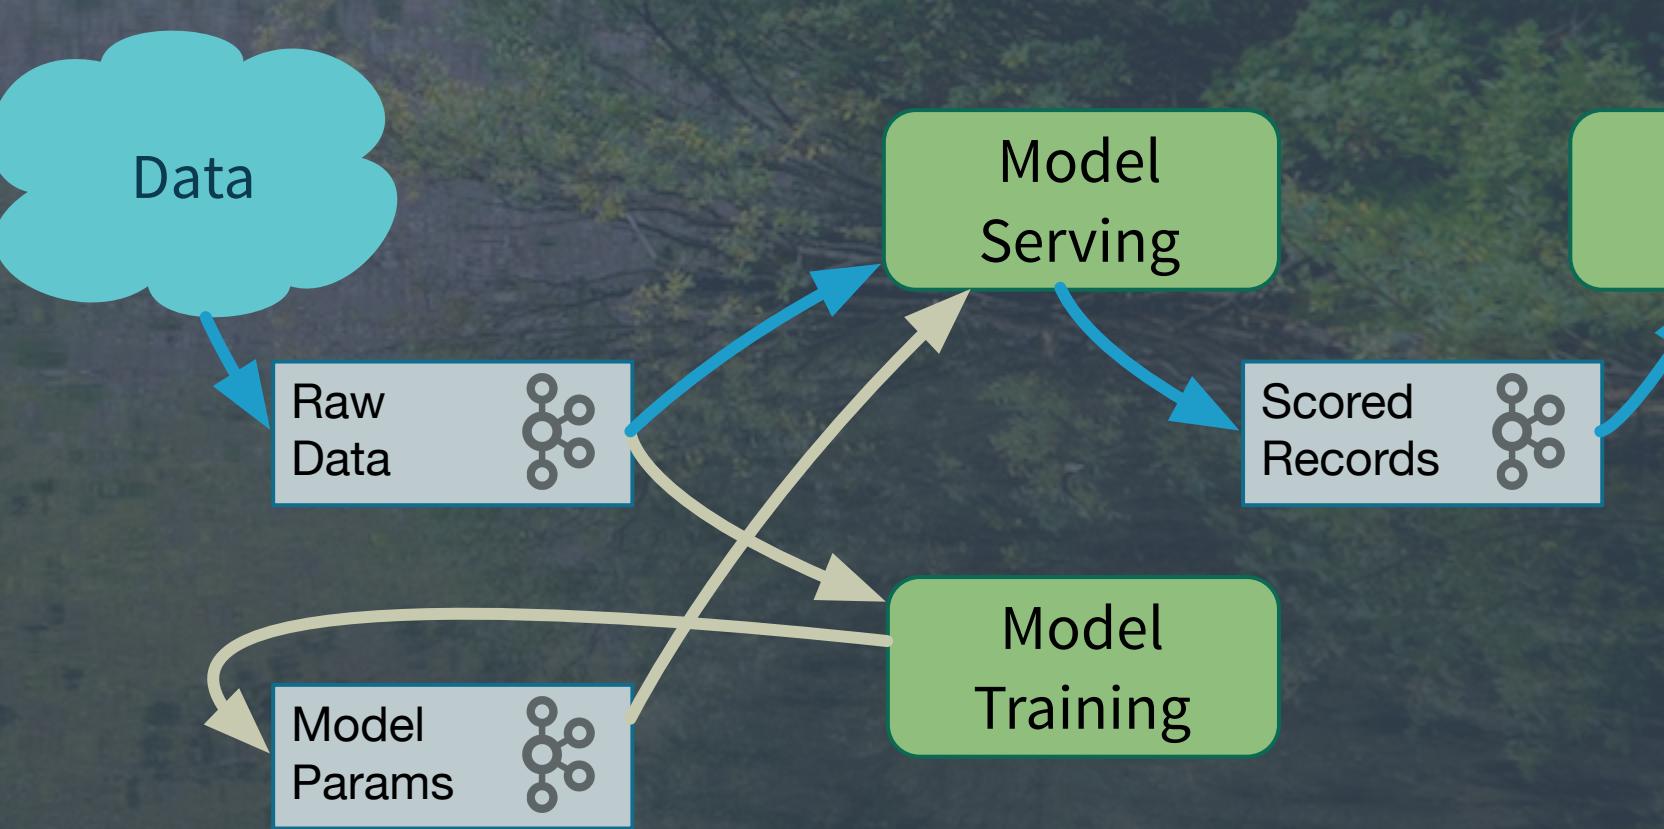


```
val builder = new StreamsBuilders // New Scala Wrapper API.
```

```
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)
```

```
val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())
```

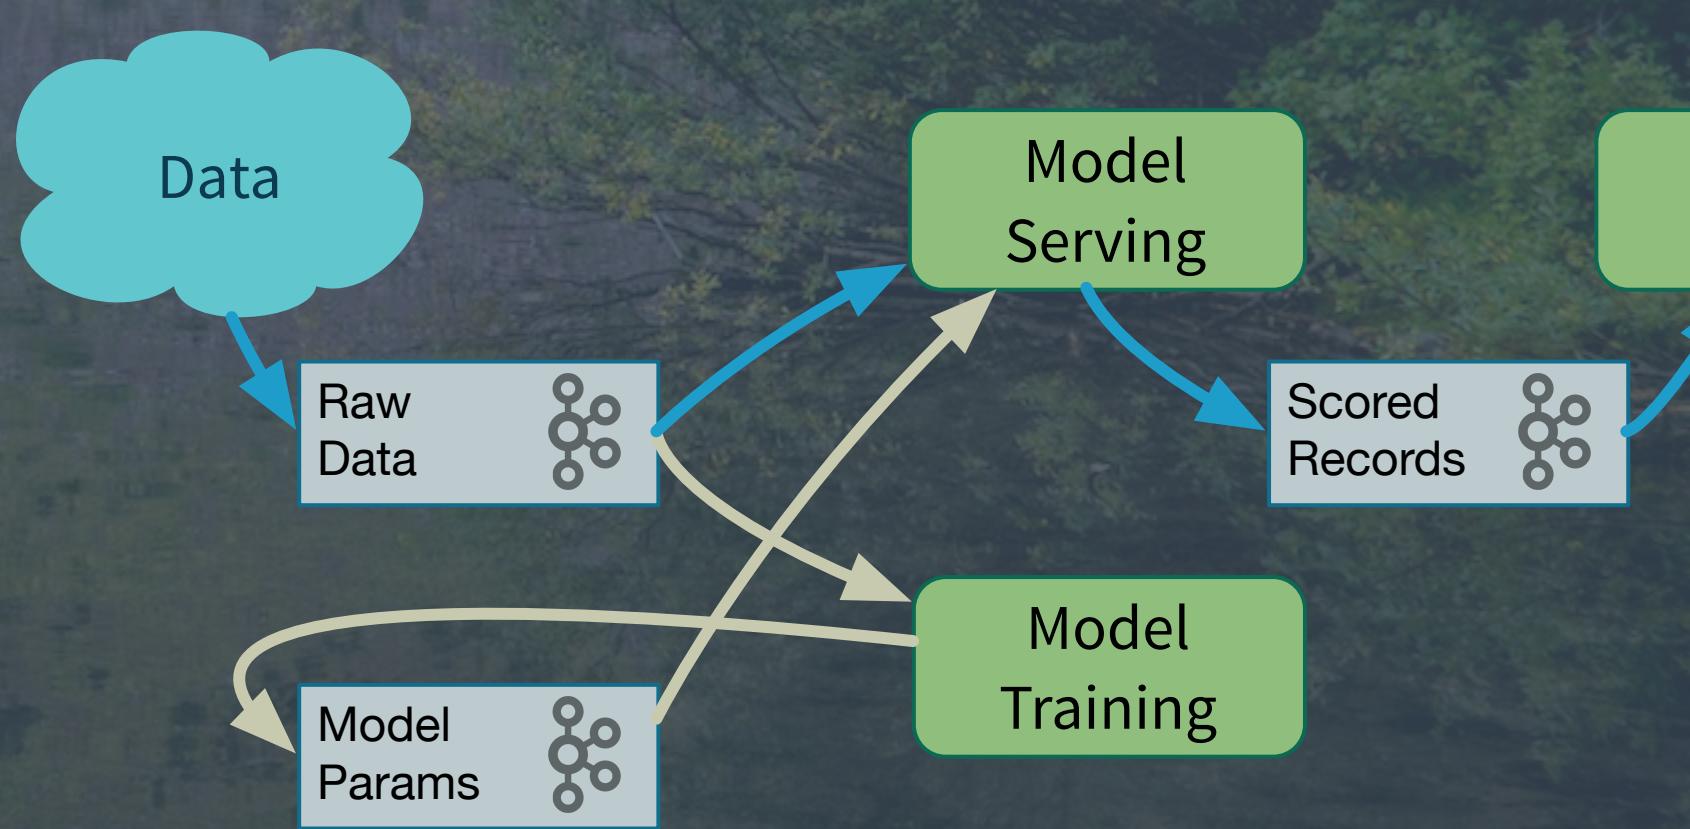


```
val builder = new StreamsBuilders // New Scala Wrapper API.
```

```
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(score(score, record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())
```



```

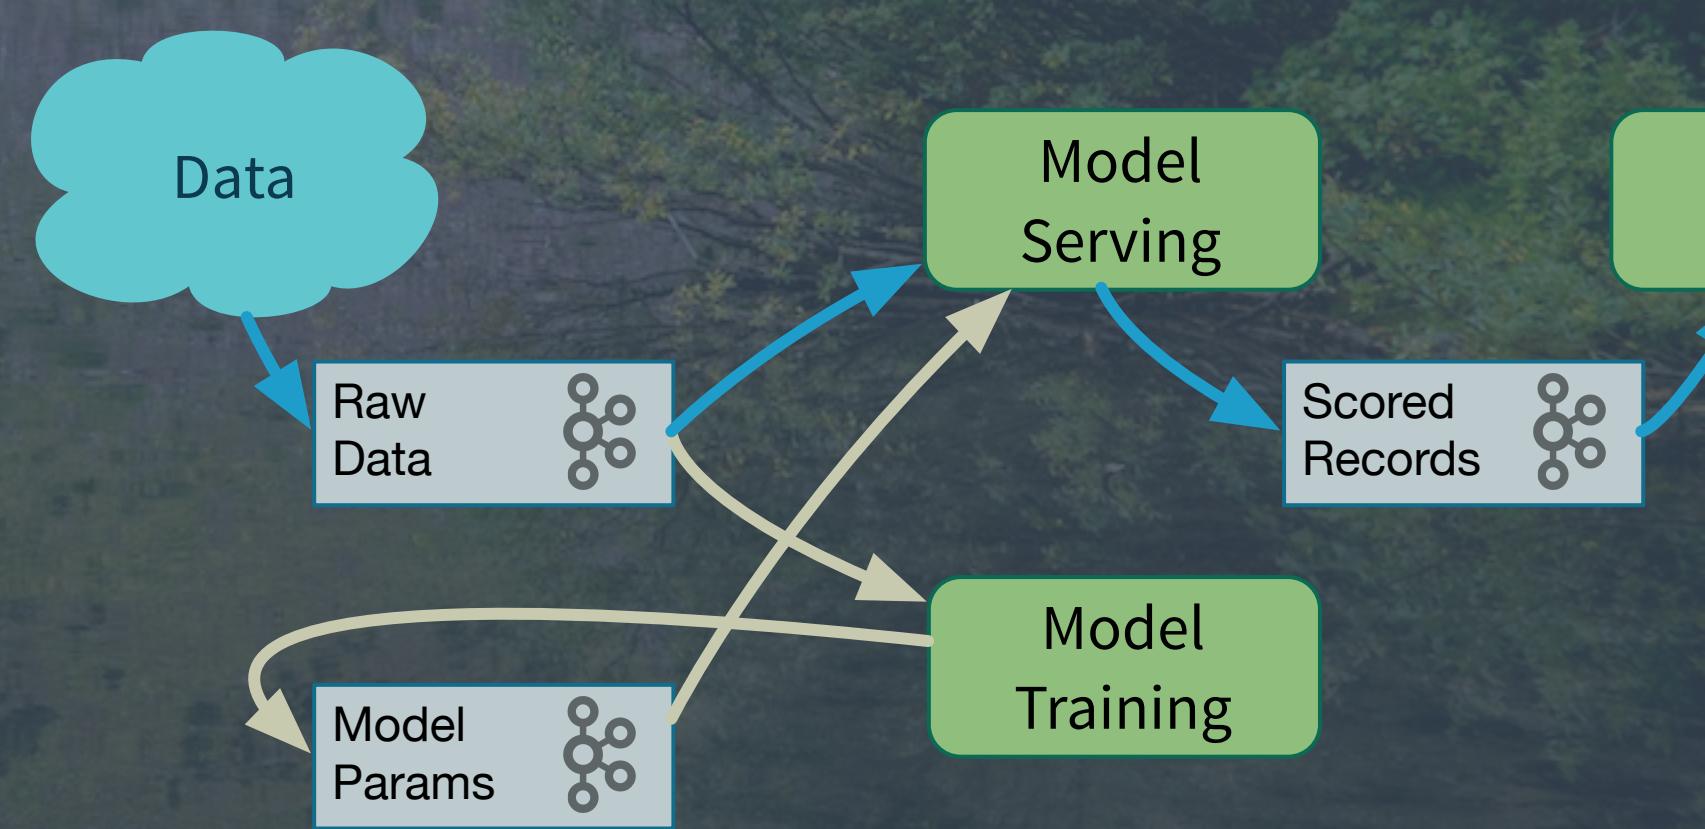
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```



```

val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

```

```

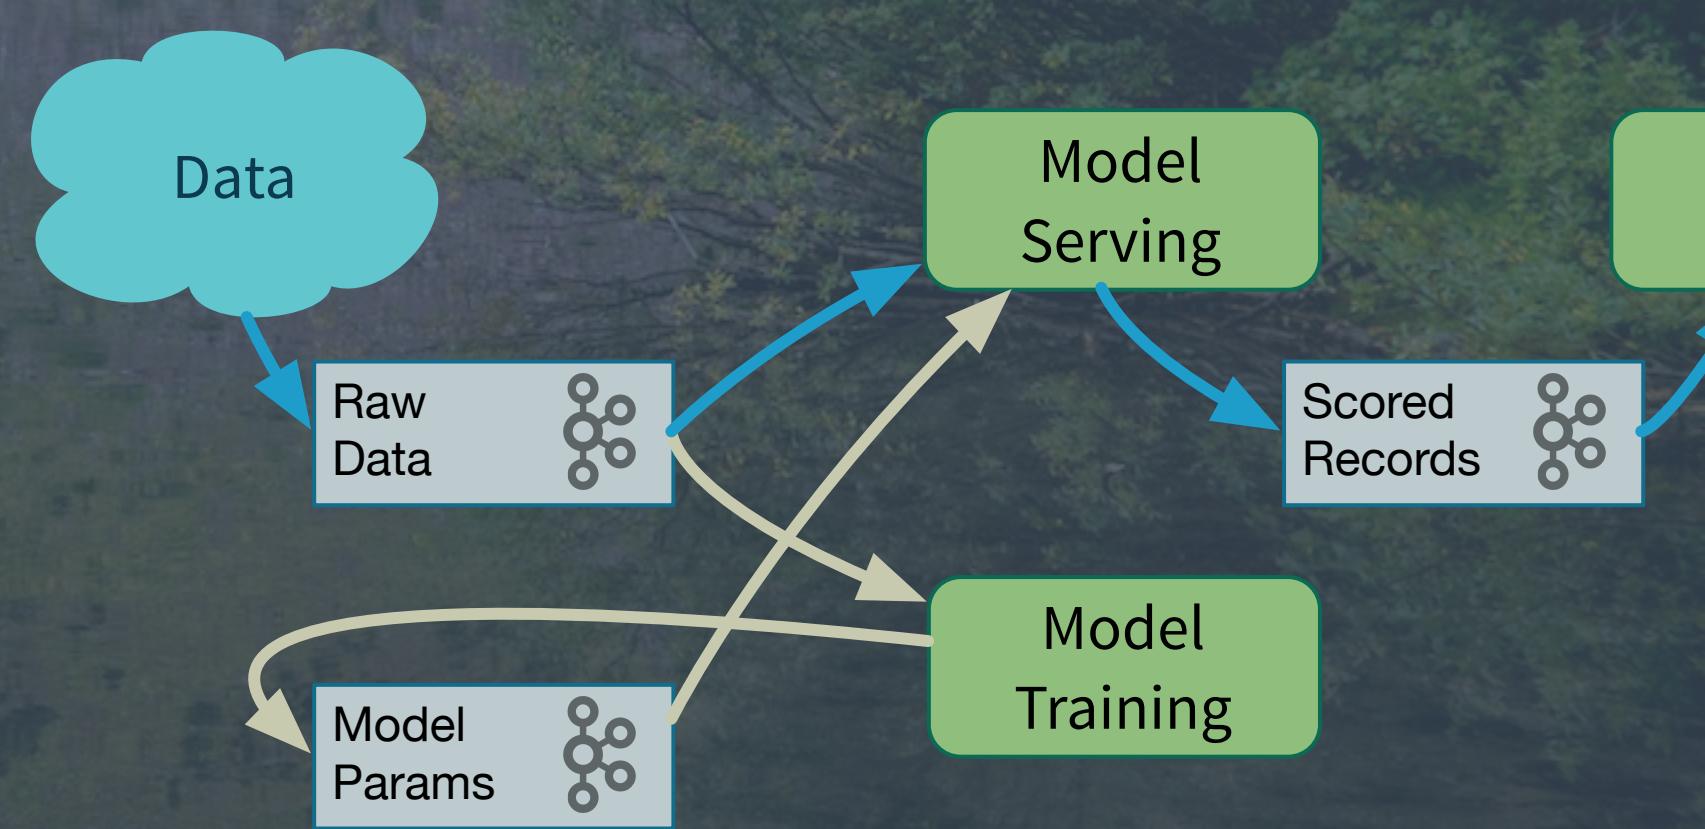
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(score(score(record), record)))
  .to(scoredRecordsTopic)

```

```

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```

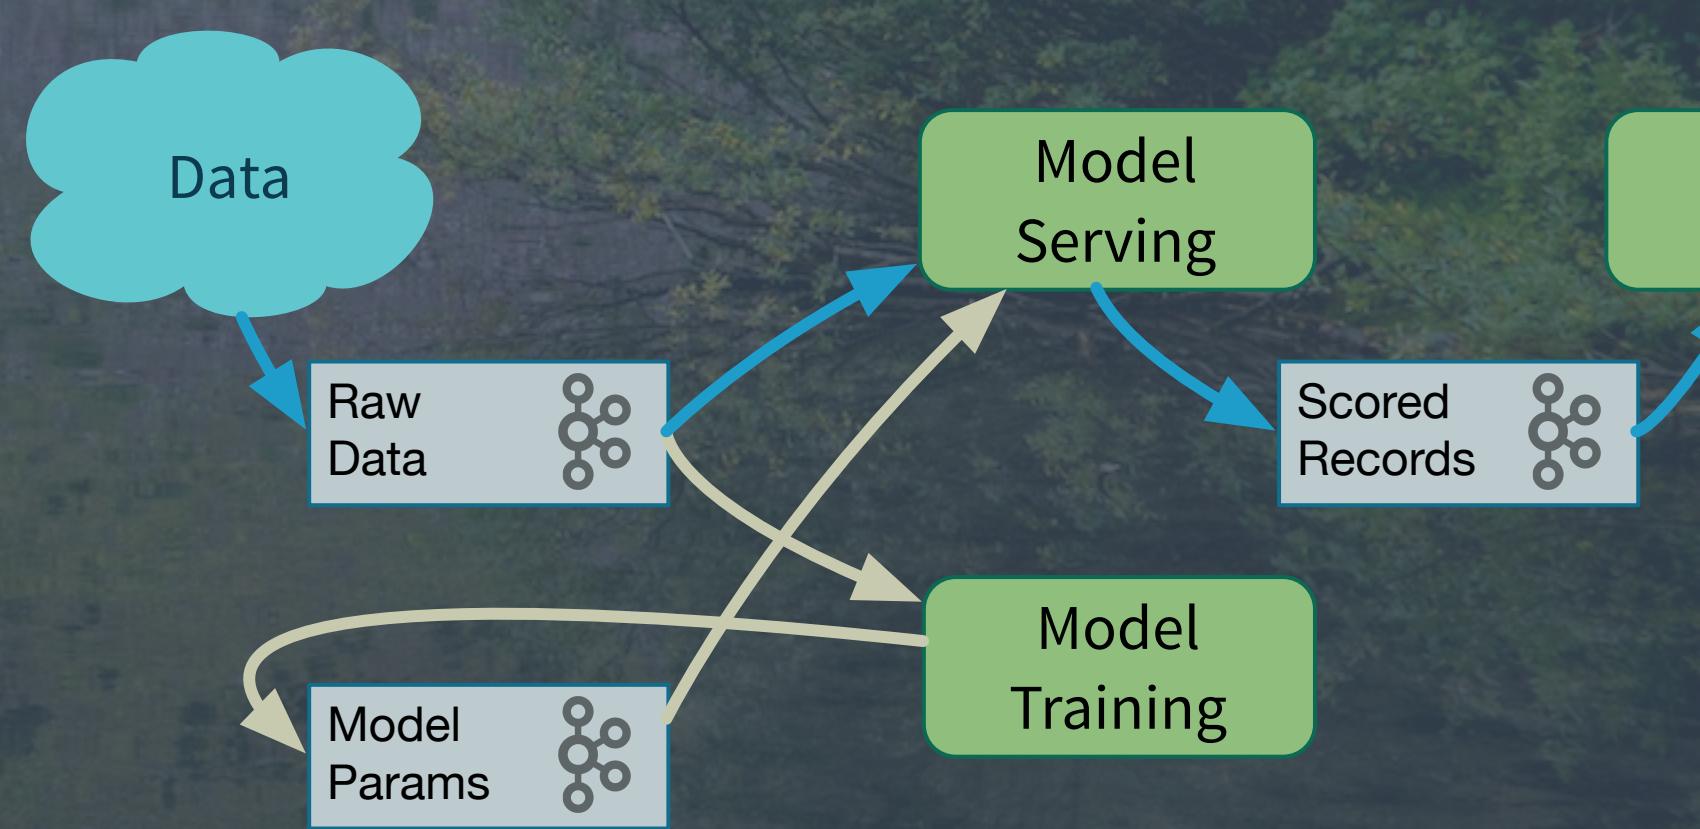


```
val builder = new StreamsBuilders // New Scala Wrapper API.
```

```
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used
```

```
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
    .filter((key, model) => model.valid) // Successful?
    .mapValues(model => ModelImpl.findModel(model))
    .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
    .filter((key, record) => record.valid)
    .mapValues(record => new ScoredRecord(scorer.score(record), record))
    .to(scoredRecordsTopic)
```

```
val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())
```



```

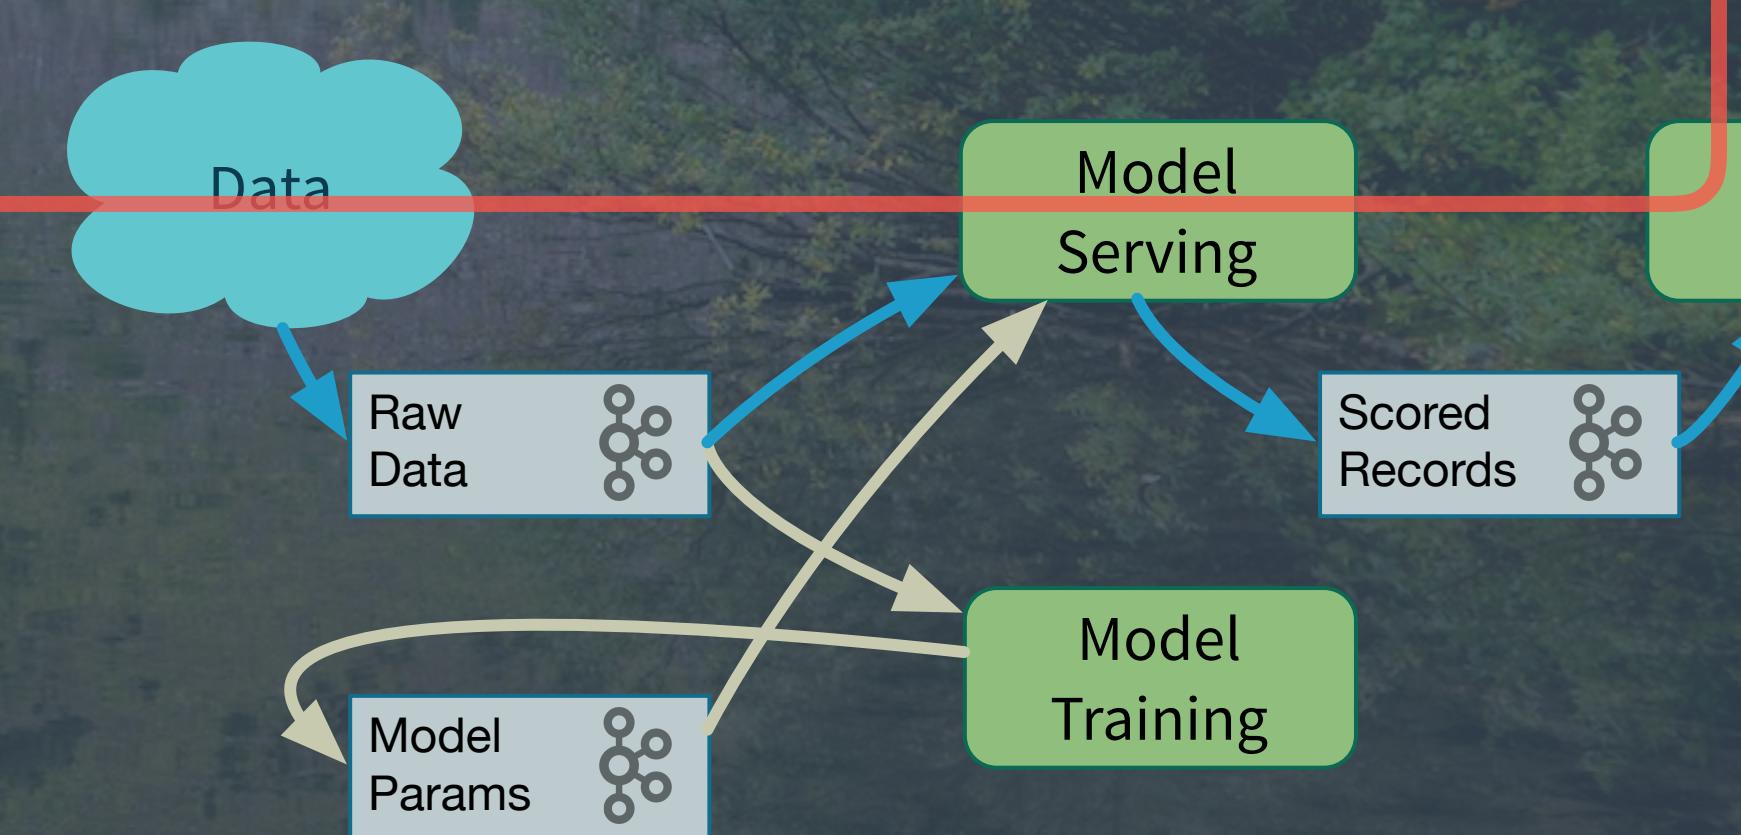
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

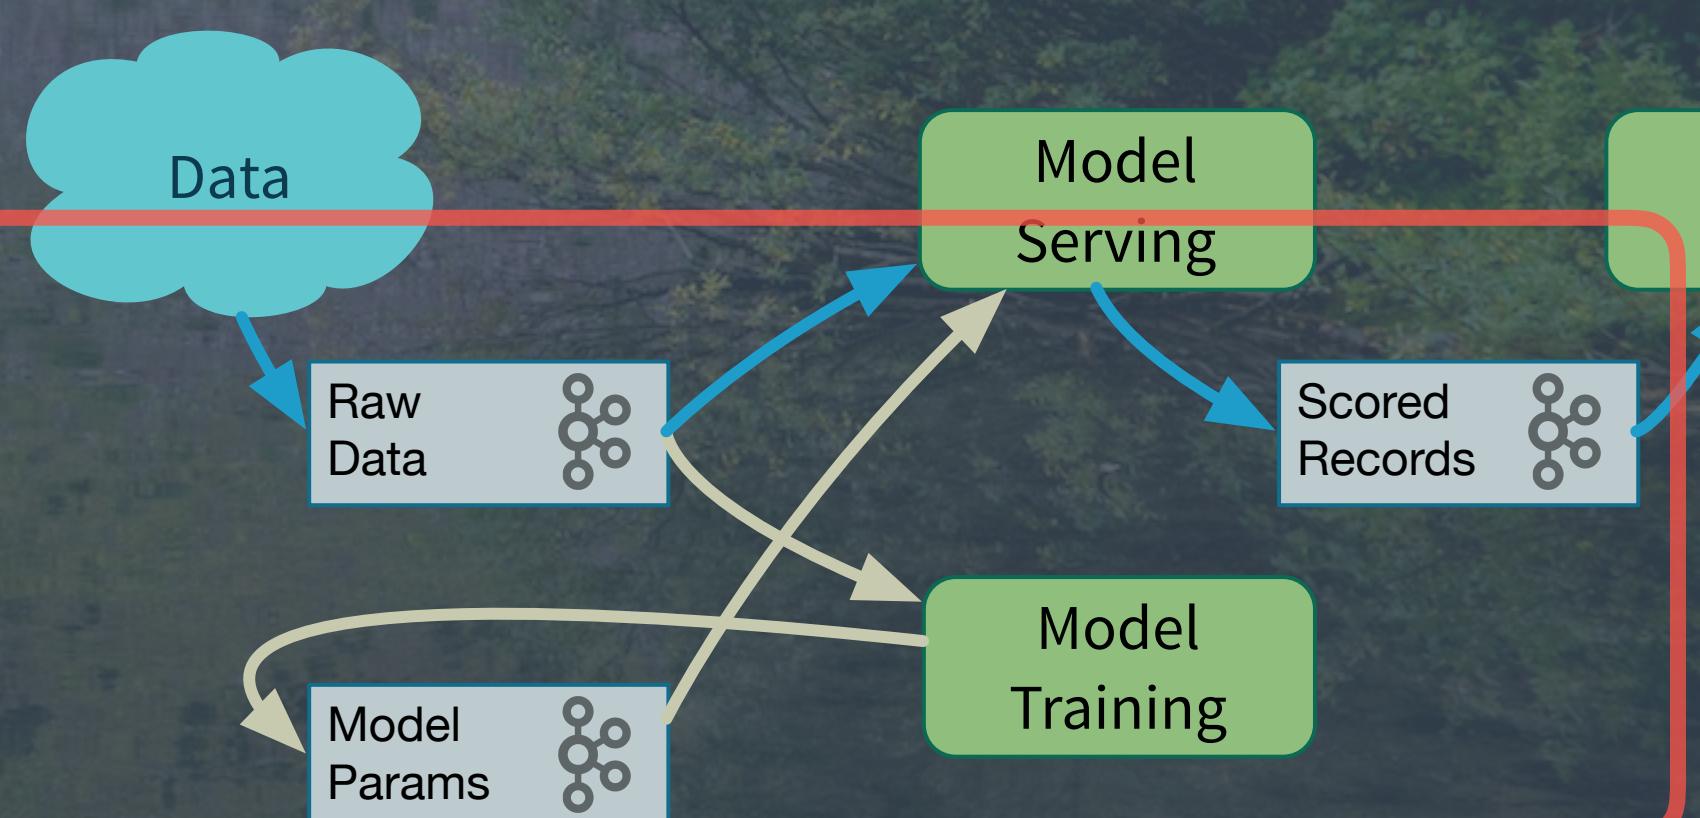
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
  data.mapValues(bytes => DataRecord.parseBytes(bytes))
    .filter((key, record) => record.valid)
    .mapValues(record => new ScoredRecord(scorer.score(record), record))
    .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```



```
val builder = new StreamsBuilders // New Scala Wrapper API.  
  
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)  
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)  
val modelProcessor = new ModelProcessor  
val scorer = new Scorer(modelProcessor) // scorer.score(record) used  
  
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record  
.filter((key, model) => model.valid) // Successful?  
.mapValues(model => ModelImpl.findModel(model))  
.process(() => modelProcessor, ...) // Set up actual model  
data.mapValues(bytes => DataRecord.parseBytes(bytes))  
.filter((key, record) => record.valid)  
.mapValues(record => new ScoredRecord(score(score(record), record)))  
.to(scoredRecordsTopic)  
  
val streams = new KafkaStreams(  
  builder.build, streamsConfiguration)  
streams.start()  
sys.addShutdownHook(streams.close())
```



What's Missing?

The rest of the microservice tools
you need...

Embed your Kafka Streams code
in microservices written with Akka
or other JVM tools.



akka Akka Streams

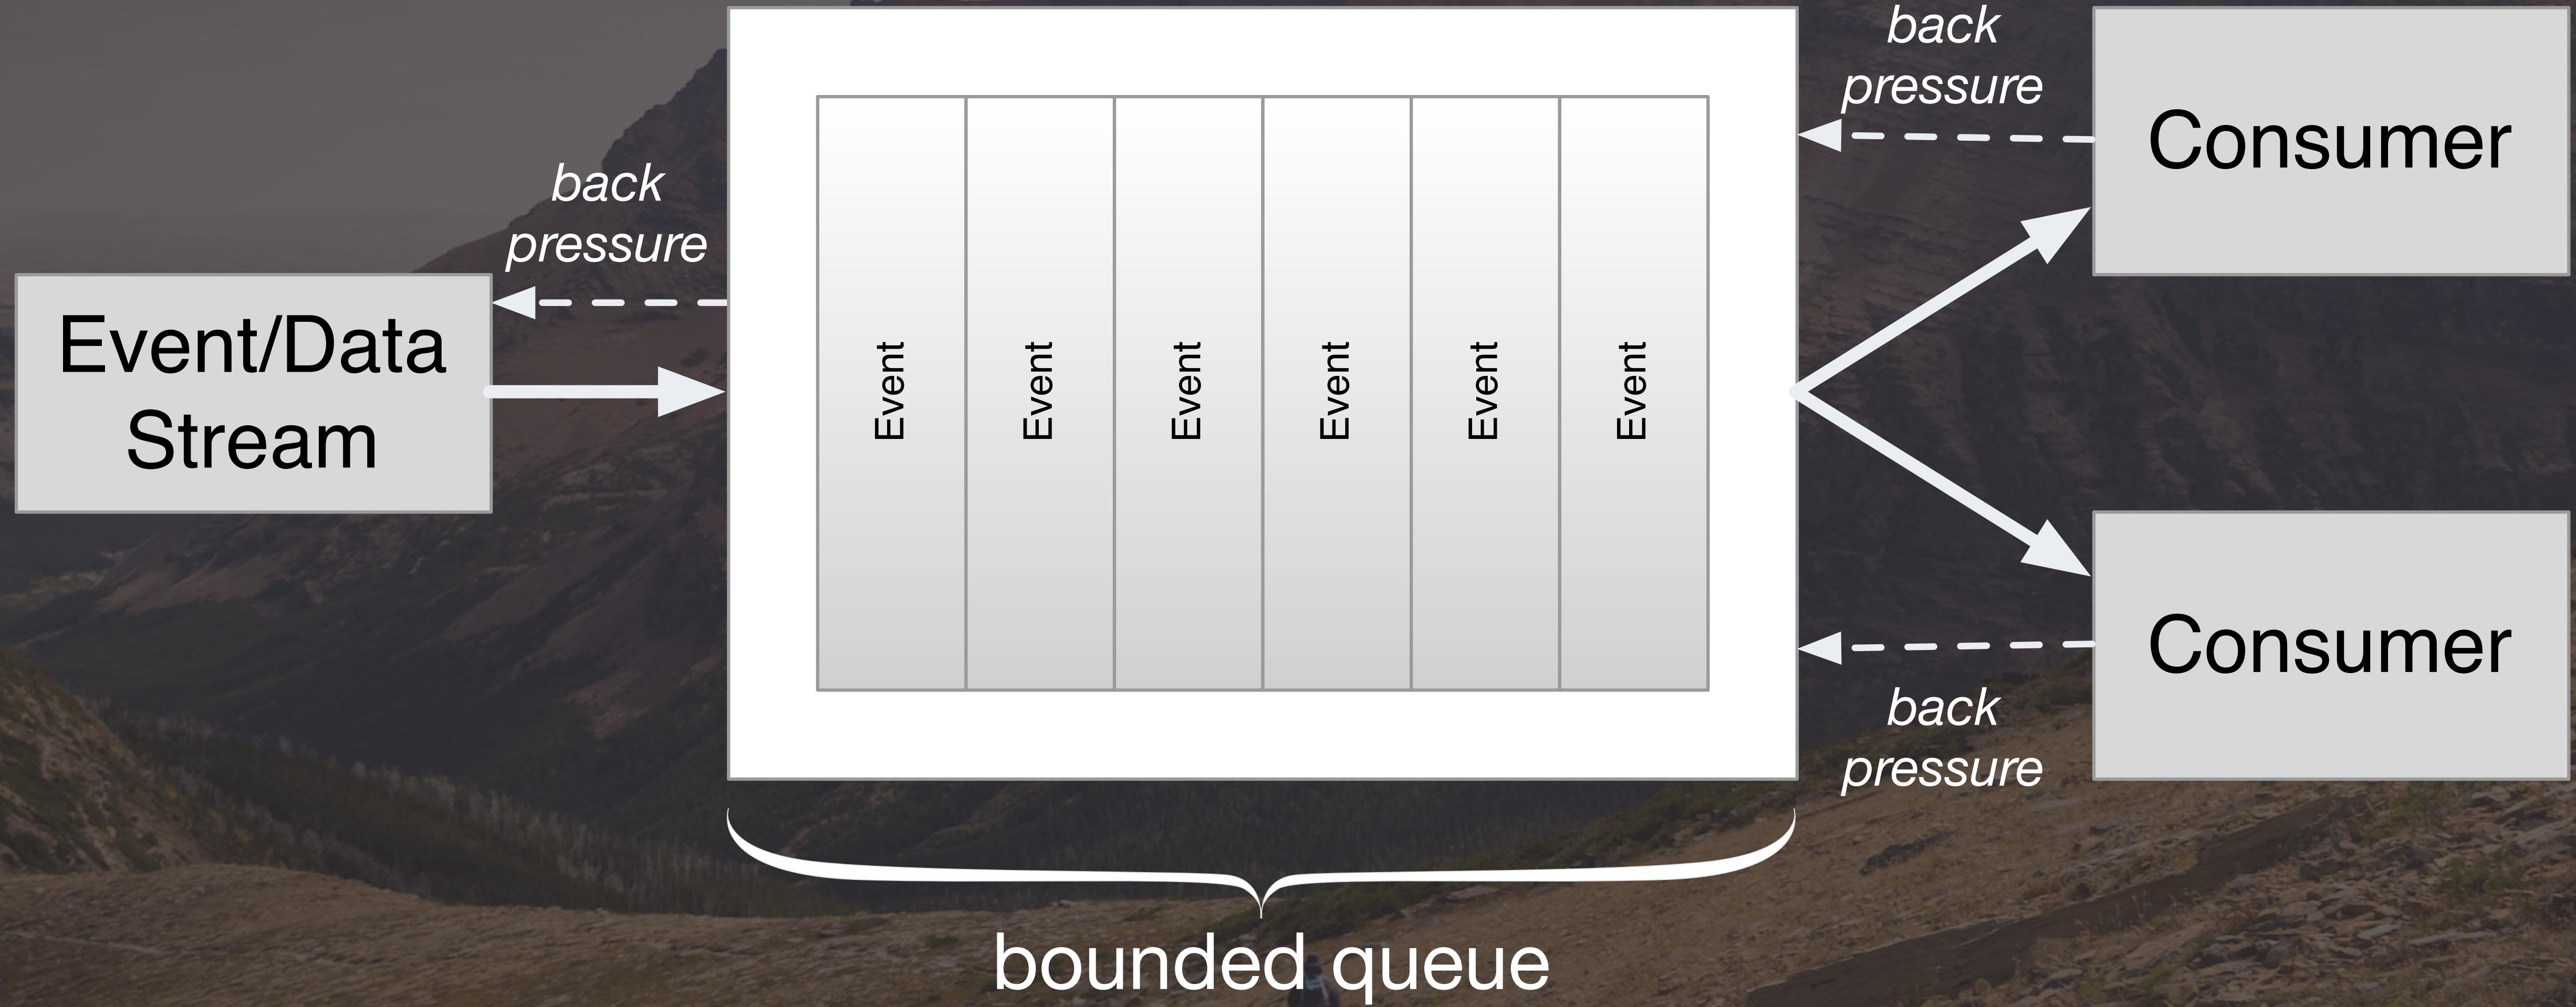


akka Akka Streams

- Implements *Reactive Streams*
 - <http://www.reactive-streams.org/>
- *Back pressure* for flow control

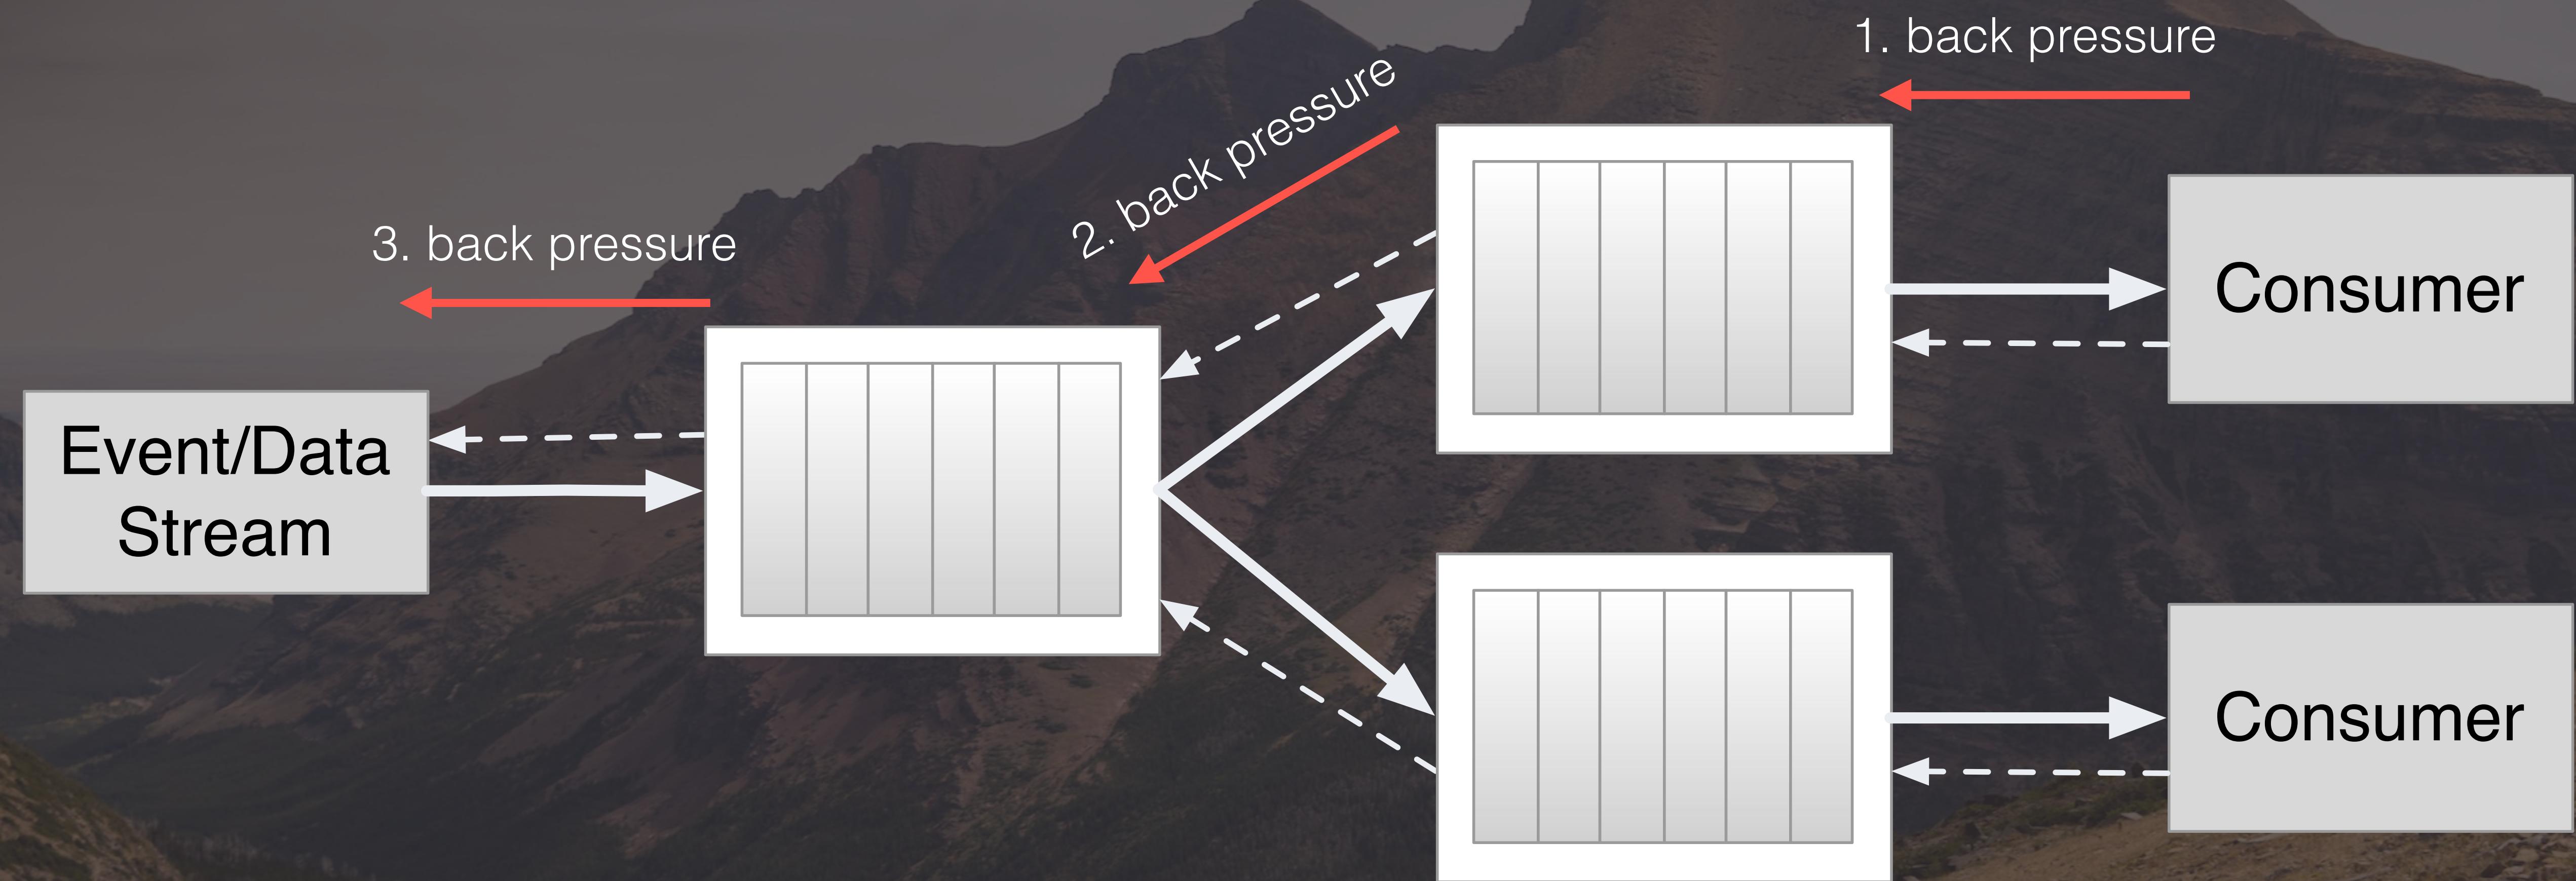


Akka Streams





Akka Streams



Back pressure composes!



akka Akka Streams

- Part of the Akka ecosystem
 - Akka Actors, Akka Cluster, Akka HTTP, Akka Persistence, ...
 - Alpakka - rich connection library
 - Optimized for low overhead and latency



Akka Streams

- The “gist” - calculate factorials:

```
val source: Source[Int, NotUsed] = Source(1 to 10)
val factorials = source.scan(BigInt(1)) (
  (total, next) => total * next )
factorials.runWith(Sink.foreach(println))
```

1
2
6
24
120
720
5040
40320
362880
3628800



Akka Streams

- The “gist” - calculate factorials:

```
val source: Source[Int, NotUsed] = Source(1 to 10)
```

```
val factorials = source.scan(BigInt(1)) (  
    (total, next) => total * next )
```

```
factorials.runWith(Sink.foreach(println))
```

A “Graph”

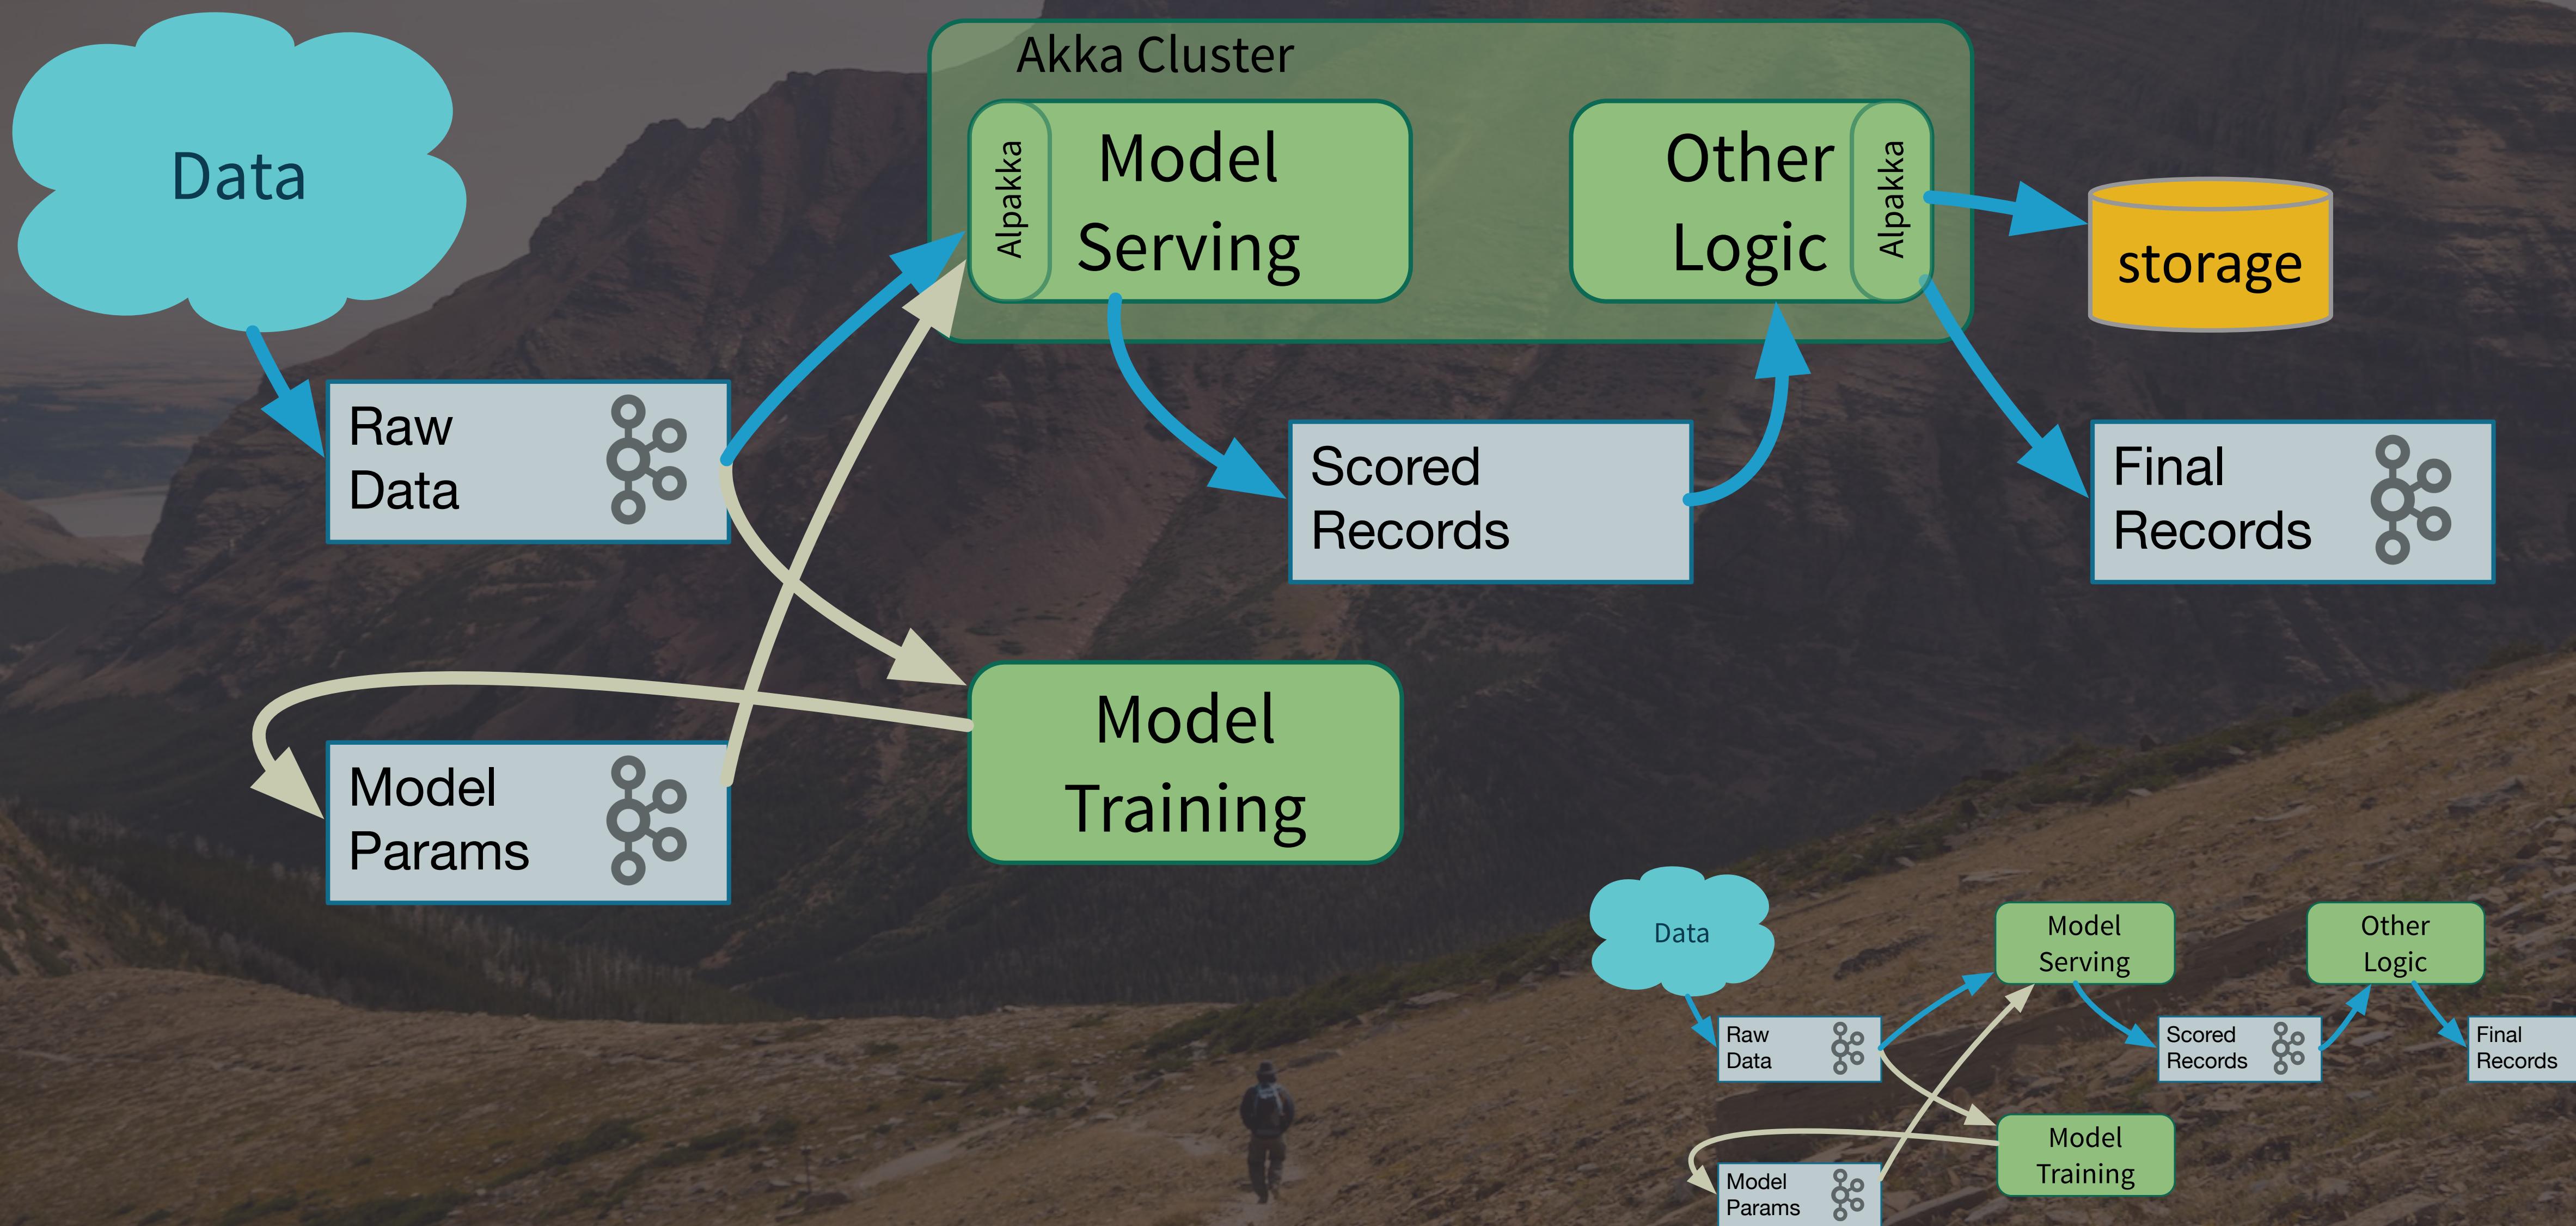
Source

Flow

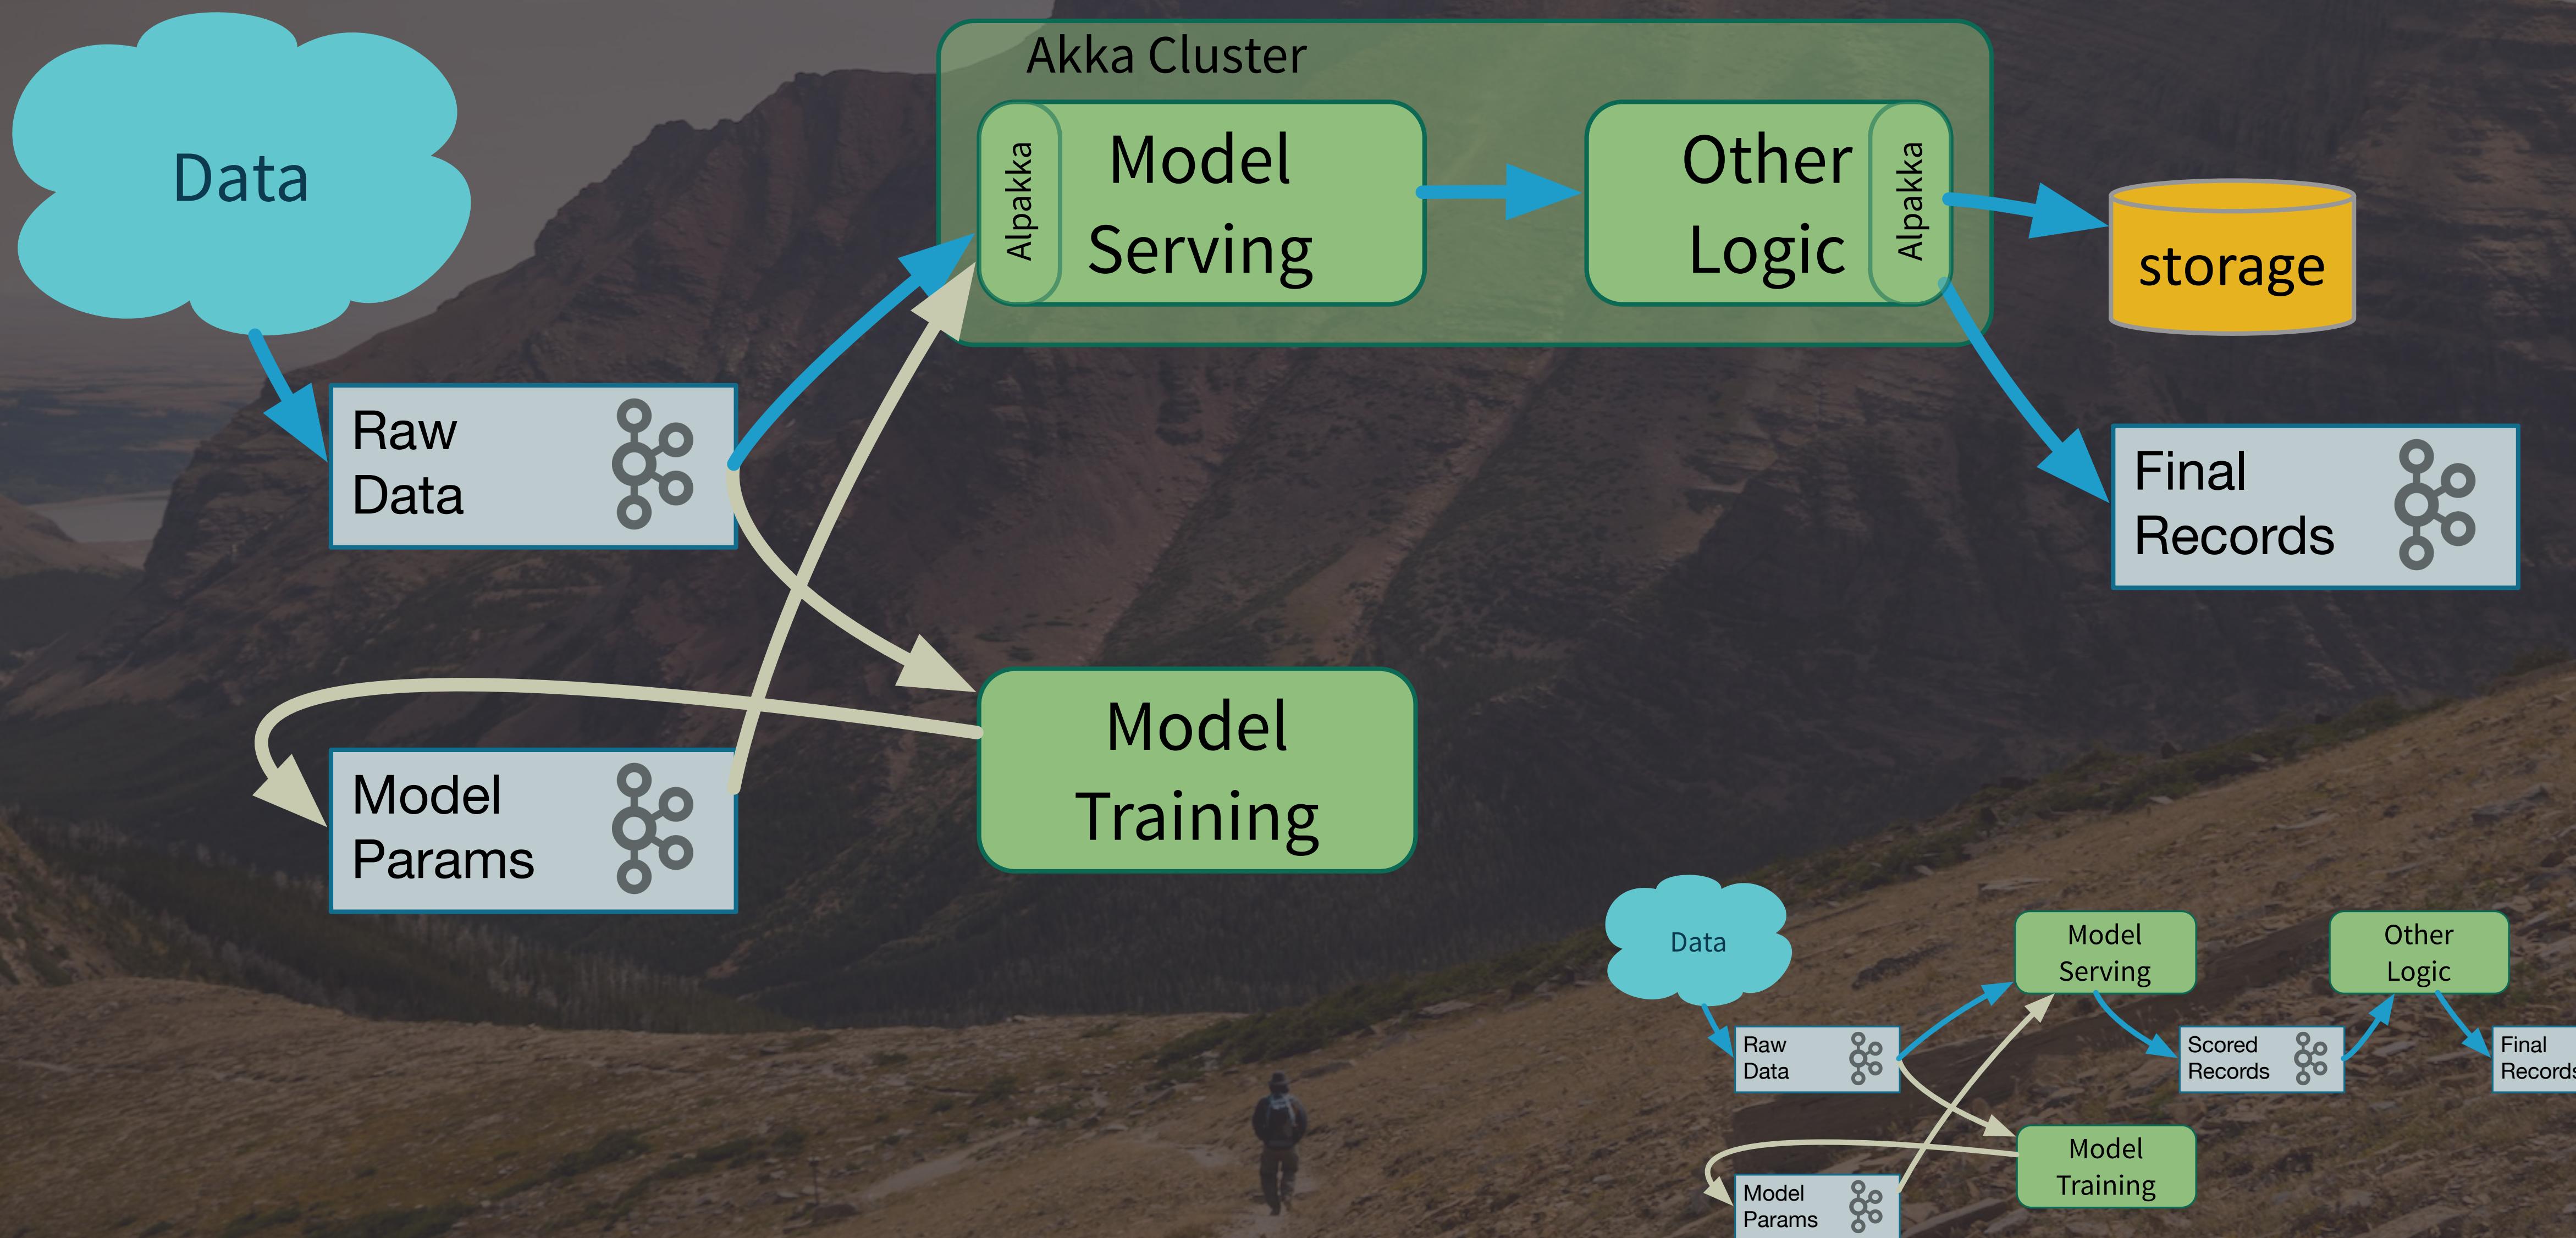
Sink

1
2
6
24
120
720
5040
40320
362880
3628800

Akka Streams Example



Akka Streams Example - “Direct” Variation

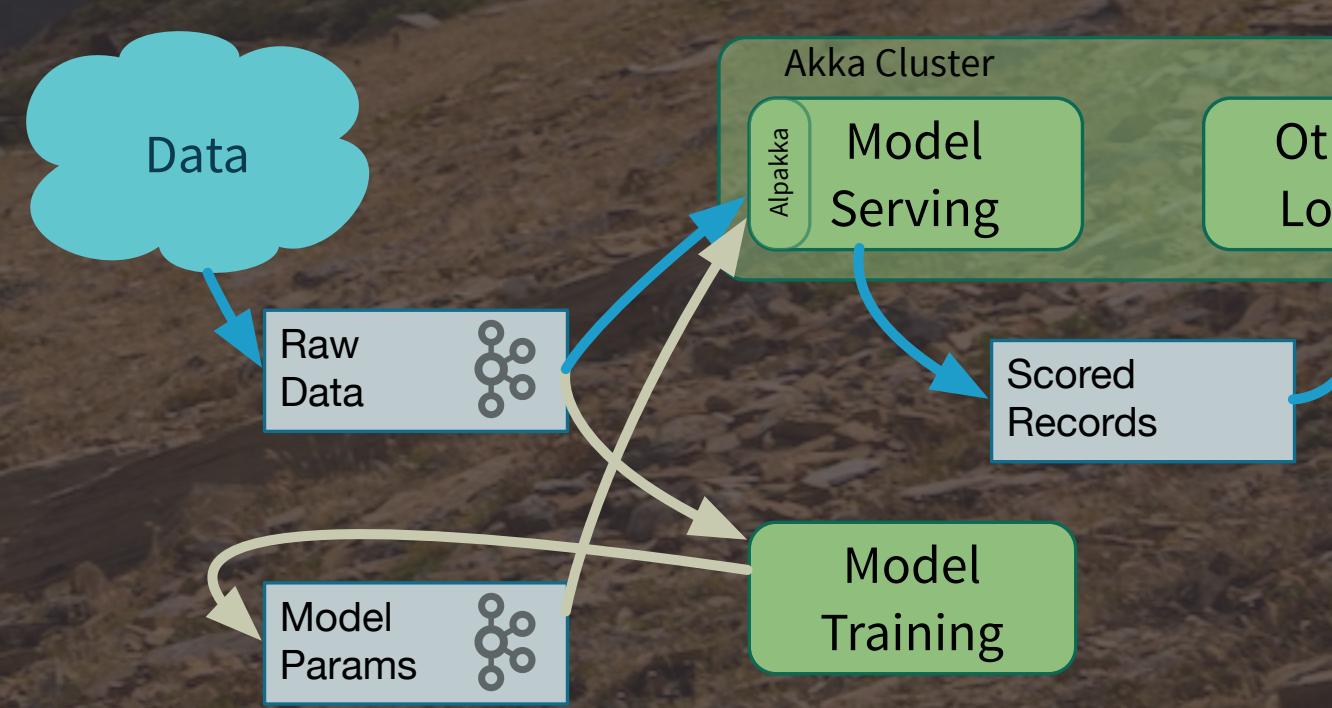


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

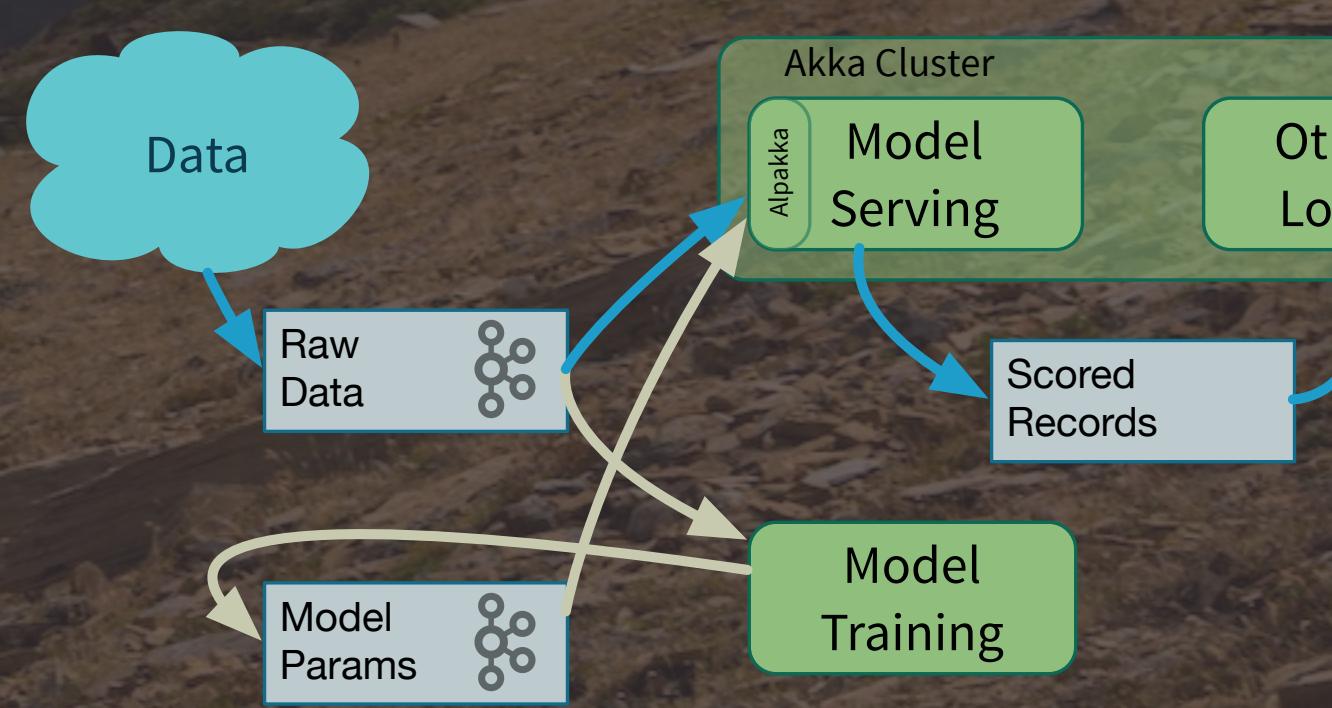


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

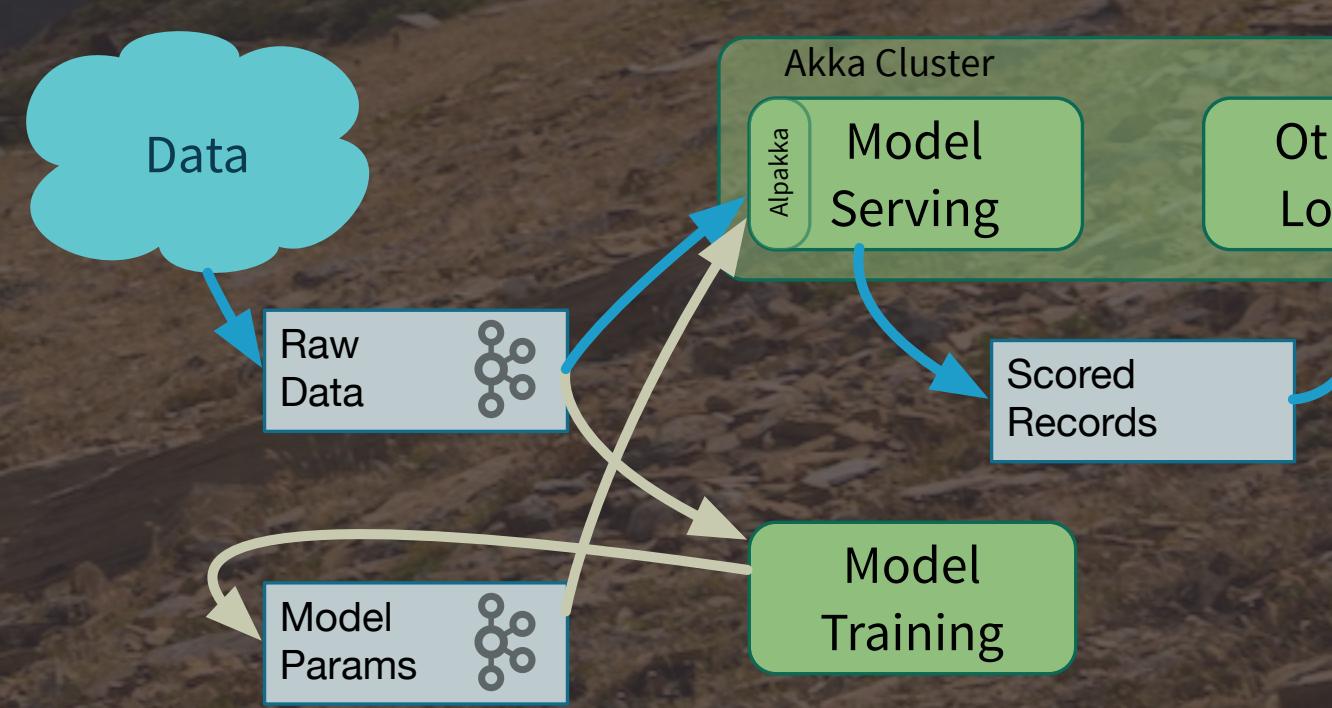


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```



```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

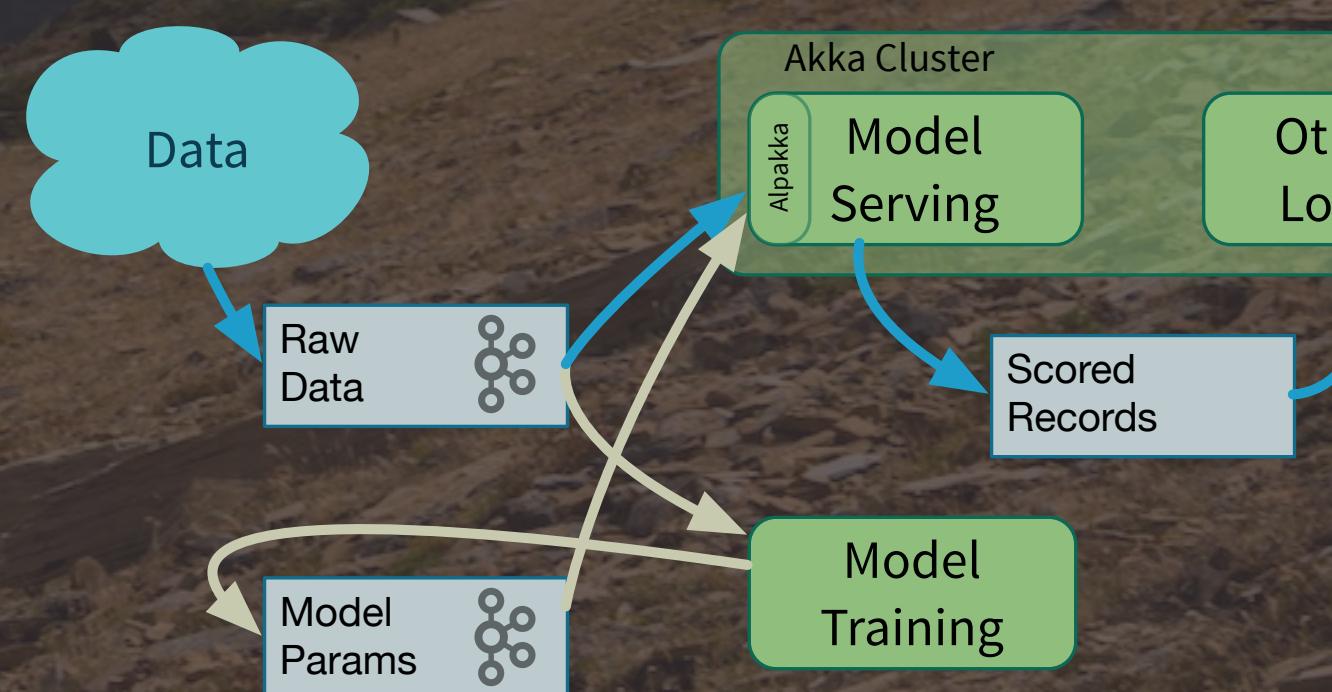
```
val modelProcessor = new ModelProcessor // Same as KS example
```

```
val scorer = new Scorer(modelProcessor) // Same as KS example
```

```
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

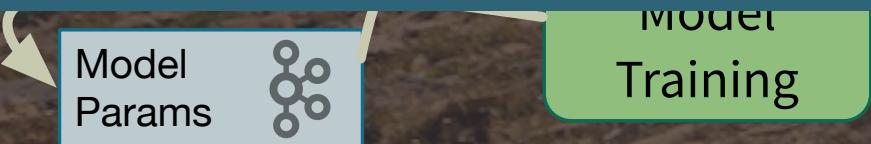
```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```



```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
import spray.actor.executionContext -> system.dispatcher
case class ModelScoringStage(scorer: ...) extends
GraphStageWithMaterializedValue[..., ...] {

    val dataRecordIn = Inlet[Record]("dataRecordIn")
    val modelRecordIn = Inlet[ModelImpl]("modelRecordIn")
    val scoringResultOut = Outlet[ScoredRecord]("scoringOut")

    ...
    setHandler(dataRecordIn, new InHandler {
        override def onPush(): Unit = {
            val record = grab(dataRecordIn)
            val newRecord = new ScoredRecord(scorer.score(record), record))
            push(scoringResultOut, Some(newRecord))
            pull(dataRecordIn)
        }
    })
    ...
}
.collect{ case Success(mod) => mod }
.map(model => ModelImpl.findModel(model))
.collect{ case Success(modImpl) => modImpl }
```

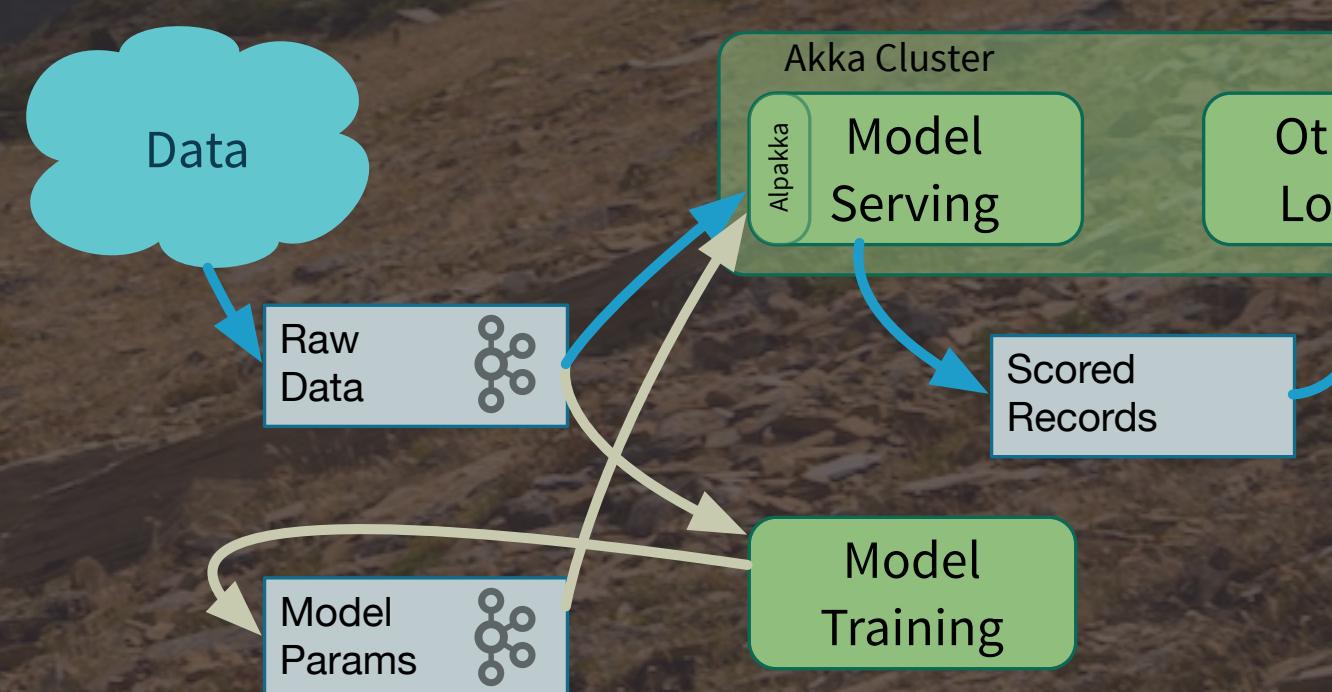


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```



```

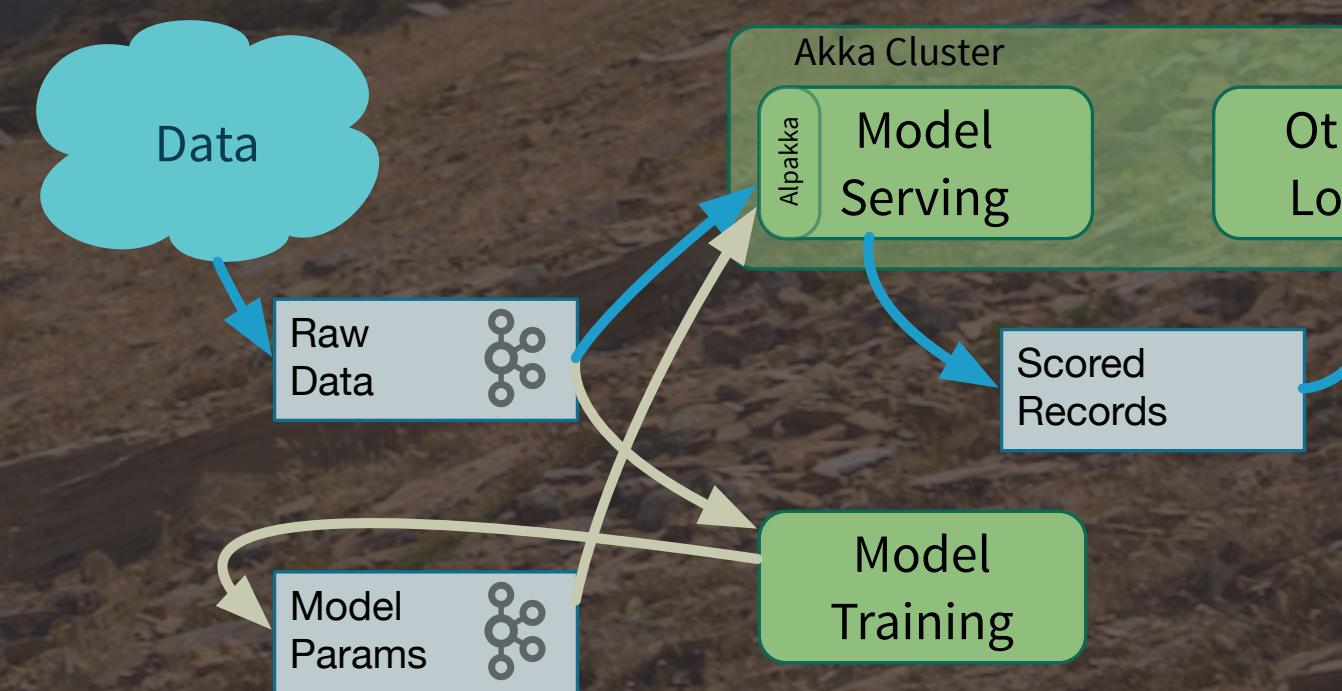
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
  .foreach(modImpl => modelProcessor.setModel(modImpl))
modelStream.to(Sink.ignore).run() // No “sinking” required; just run

```

```

dataStream
  .viaMat(modelScoringStage)(Keep.right)
  .map(result => new ProducerRecord[Array[Byte], ScoredRecord](
    scoredRecordsTopic, result))
  .runWith(Producer.plainSink(producerSettings))

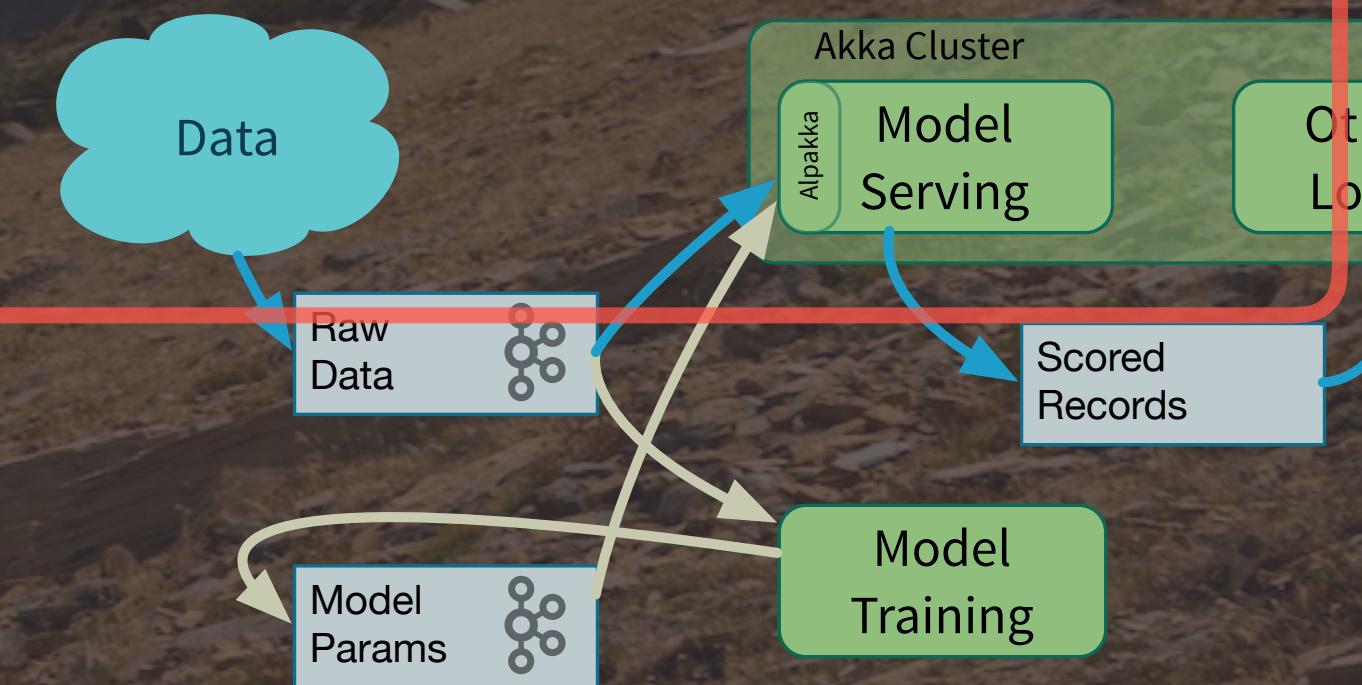
```



```
.collect{ case Success(data) -> data }

val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
  .foreach(modImpl => modelProcessor.setModel(modImpl))
modelStream.to(Sink.ignore).run() // No “sinking” required; just run
```

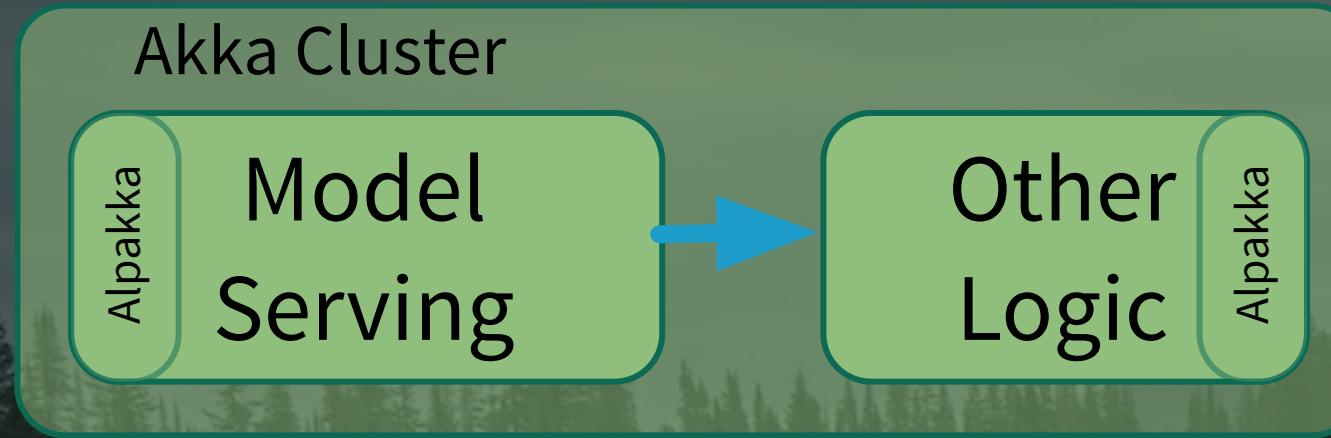
```
dataStream
  .viaMat(modelScoringStage)(Keep.right)
  .map(result => new ProducerRecord[Array[Byte], ScoredRecord](
    scoredRecordsTopic, result))
  .runWith(Producer.plainSink(producerSettings))
```



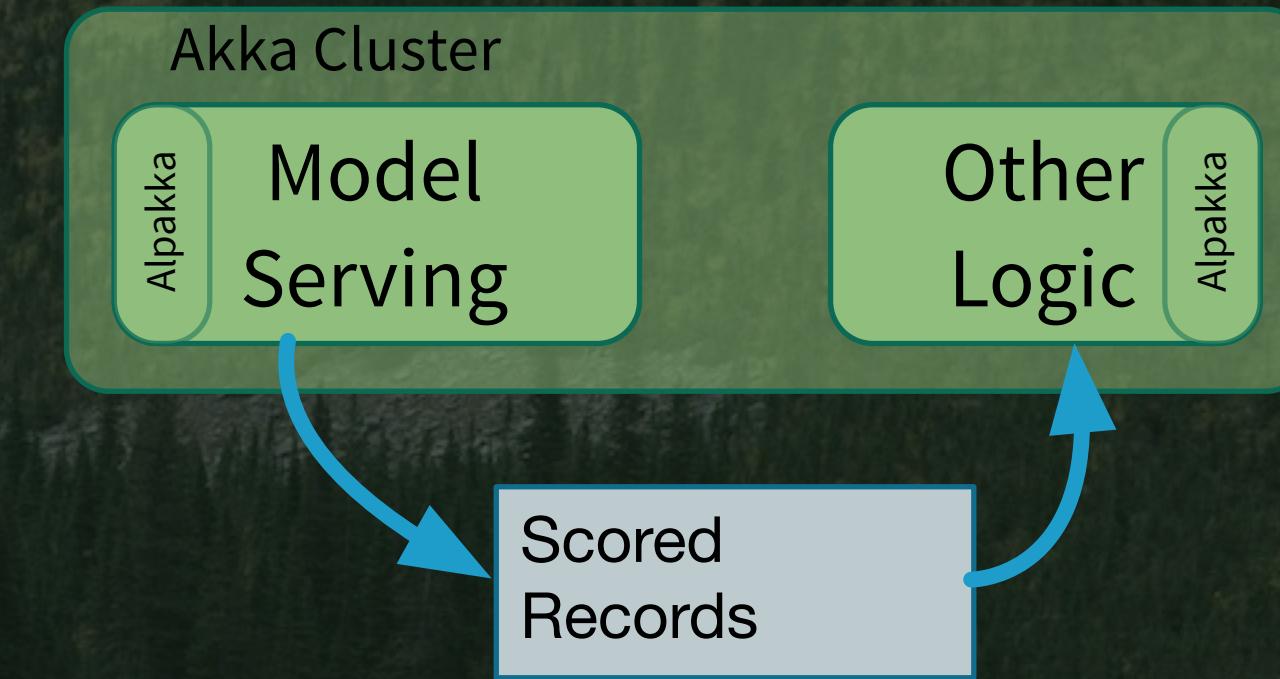
A wide-angle photograph of a serene mountain lake. The water is still, creating a perfect mirror for the surrounding environment. On the left, a dense forest of tall evergreen trees stands along the shoreline. To the right, a massive, steep mountain slope covered in green coniferous forests rises sharply. The sky above is filled with heavy, grey clouds, with some lighter areas suggesting a setting or rising sun.

Other
Concerns

Go Direct or Through Kafka?

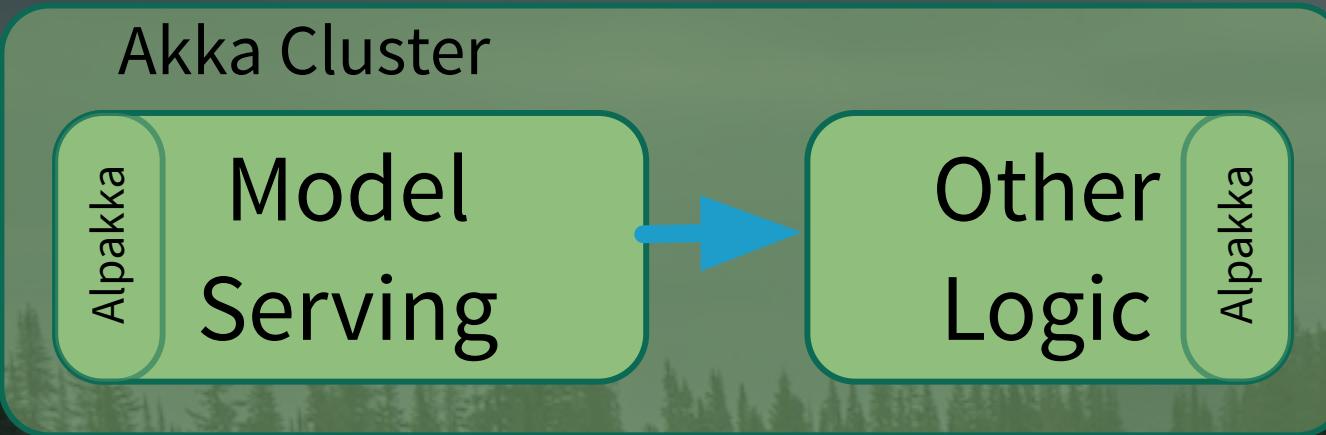


VS.

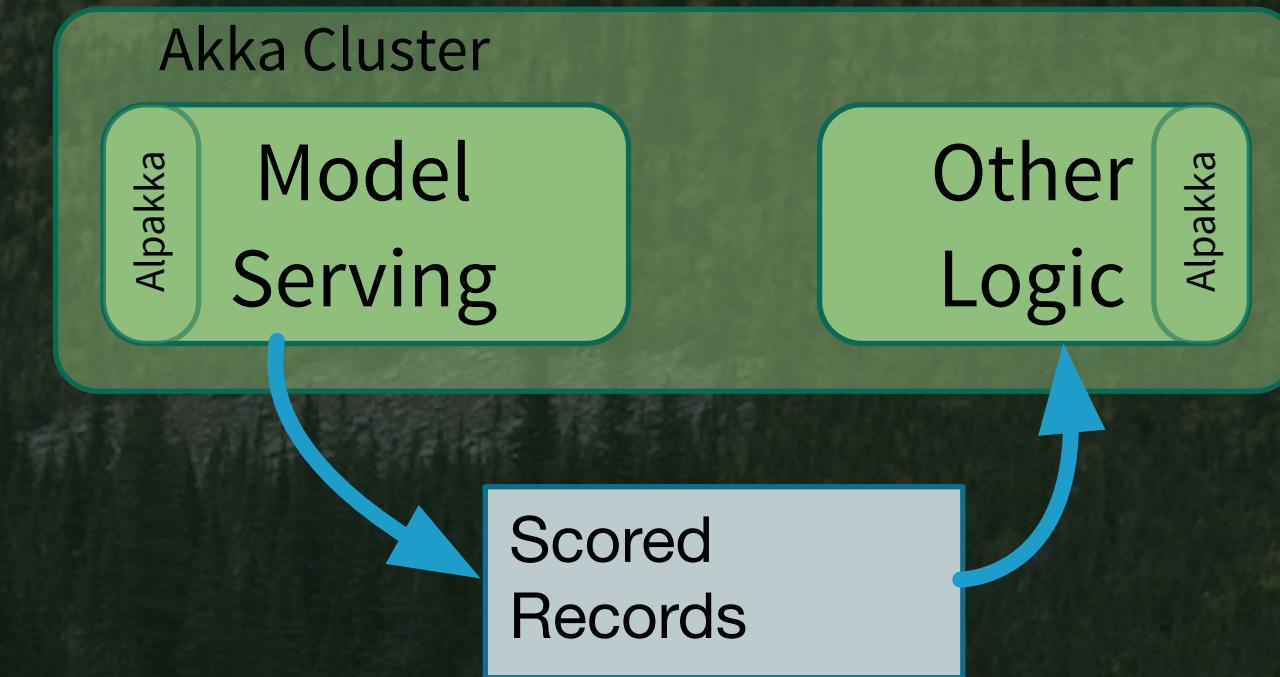


- Extremely low latency
- Higher latency (network, queue depth)

Go Direct or Through Kafka?

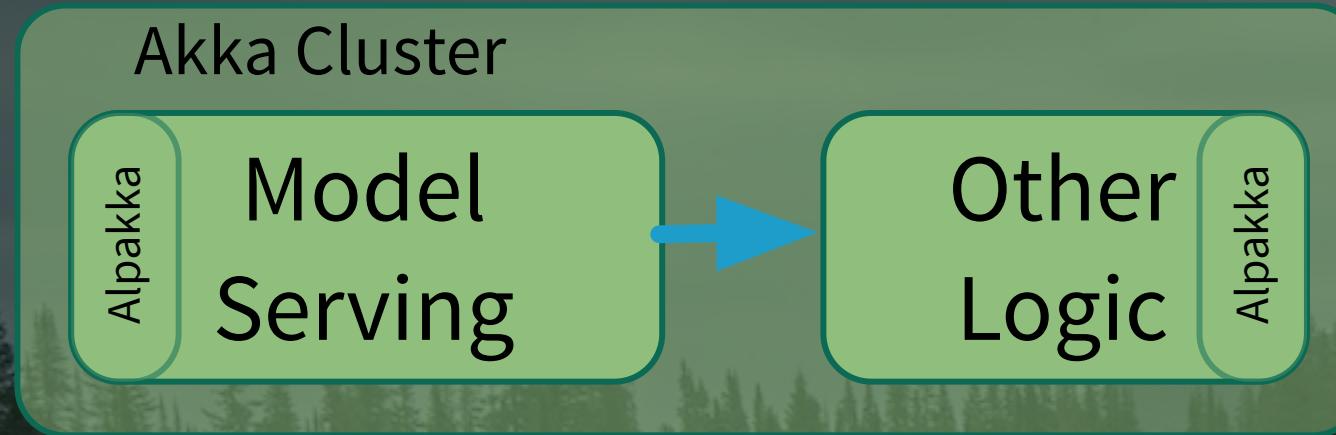


VS.

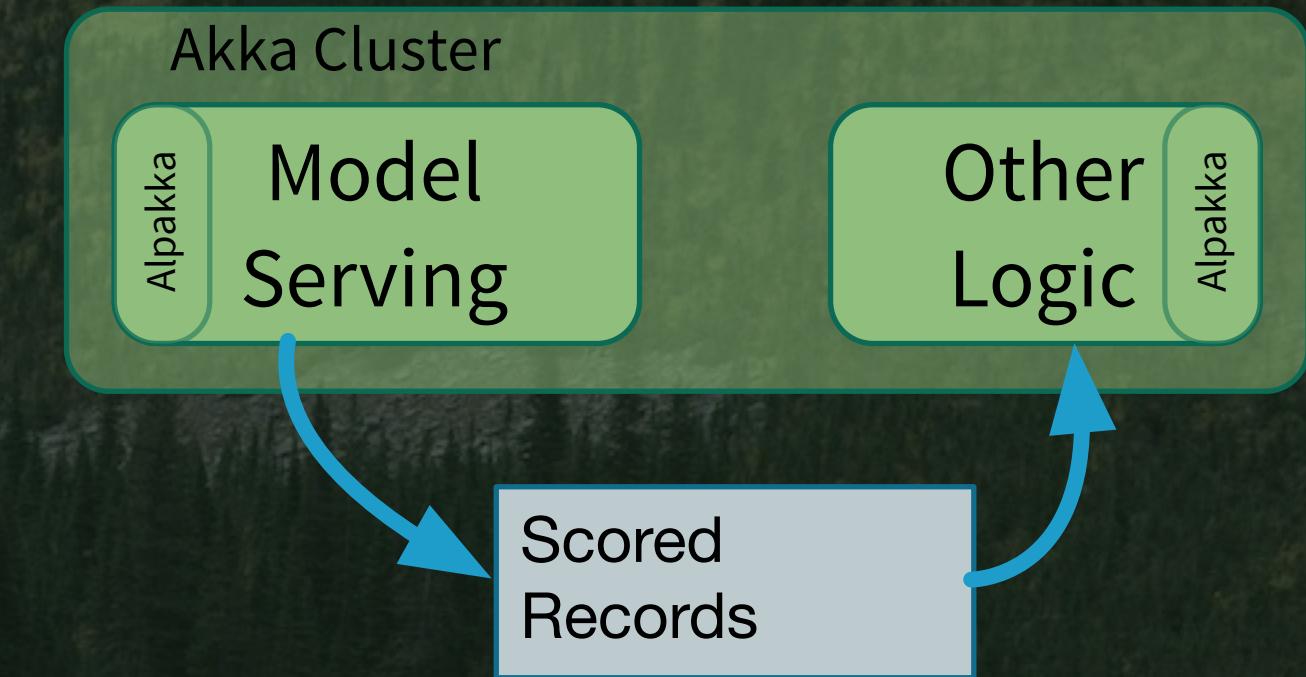


- Extremely low latency
- Minimal I/O and memory overhead. No marshaling overhead
- Higher latency (network, queue depth)
- Higher I/O and processing (marshaling) overhead

Go Direct or Through Kafka?

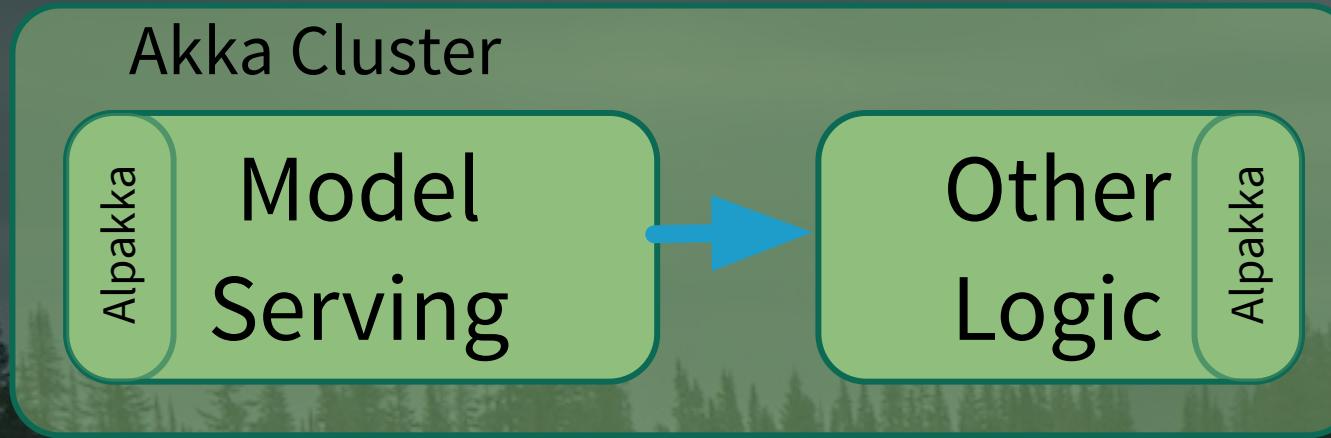


VS.

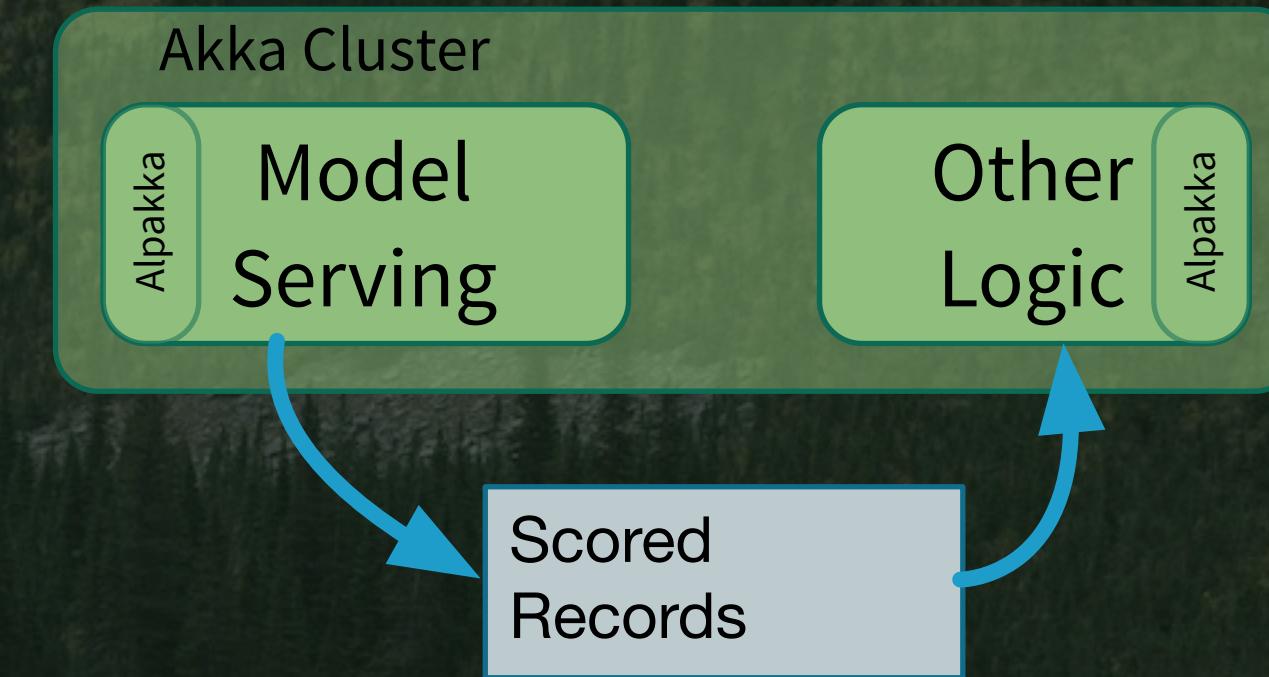


- Extremely low latency
- Minimal I/O and memory overhead. No marshaling overhead
- Hard to scale, evolve independently
- Higher latency (network, queue depth)
- Higher I/O and processing (marshaling) overhead
- Easy independent scalability, evolution

Go Direct or Through Kafka?

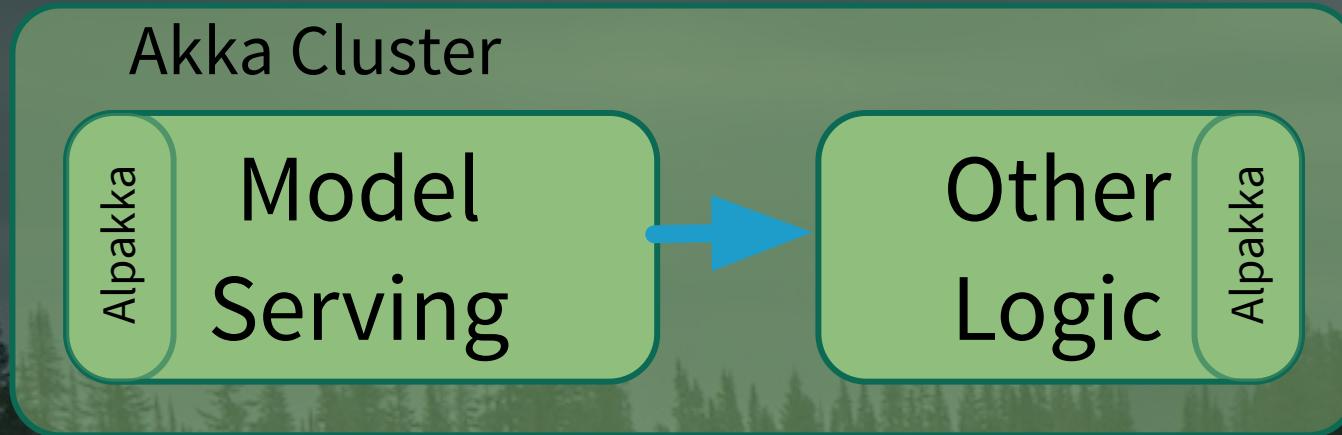


VS.

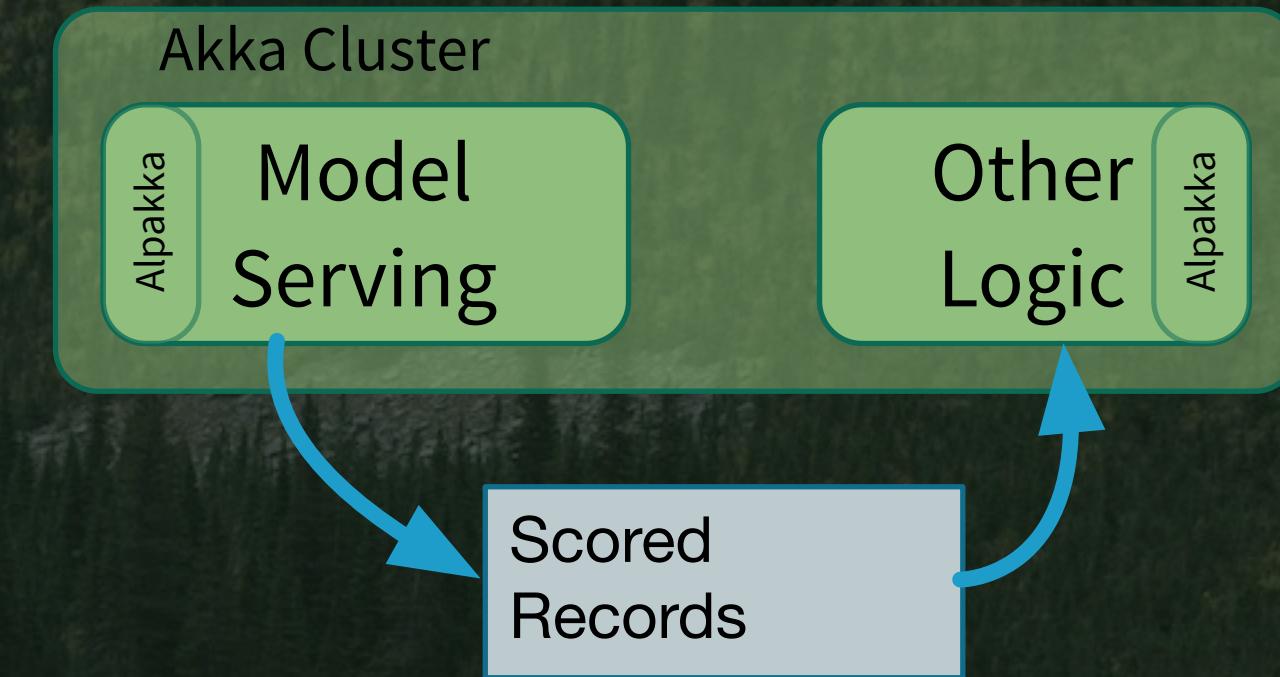


- *Reactive Streams* back pressure
- Very deep buffer (partition limited by disk size)

Go Direct or Through Kafka?



VS.



- *Reactive Streams* back pressure
- “Direct” coupling between sender and receiver, but actually uses a URL abstraction

- Very deep buffer (partition limited by disk size)
- Strong decoupling - M producers, N consumers, completely disconnected

Wrapping Up...



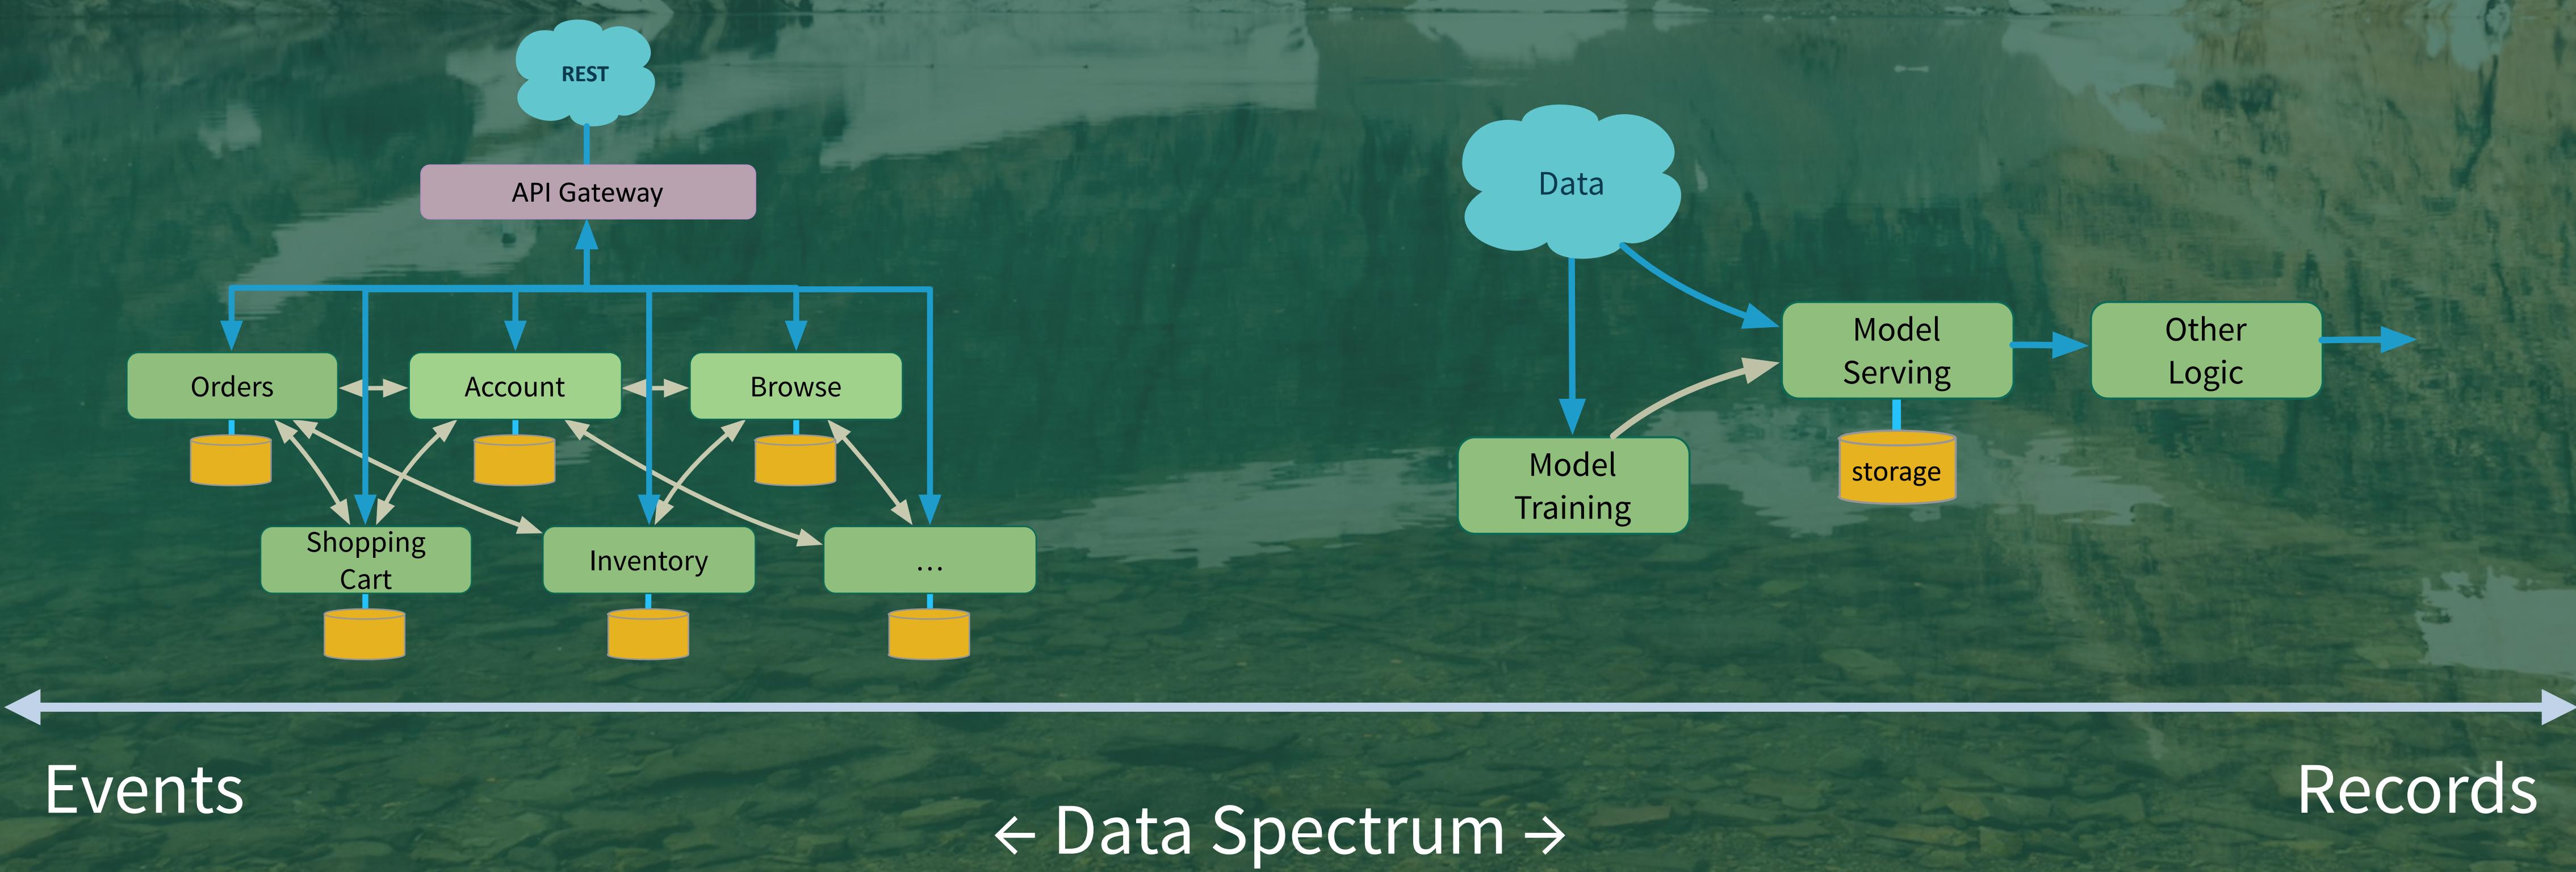
A Spectrum of Microservices

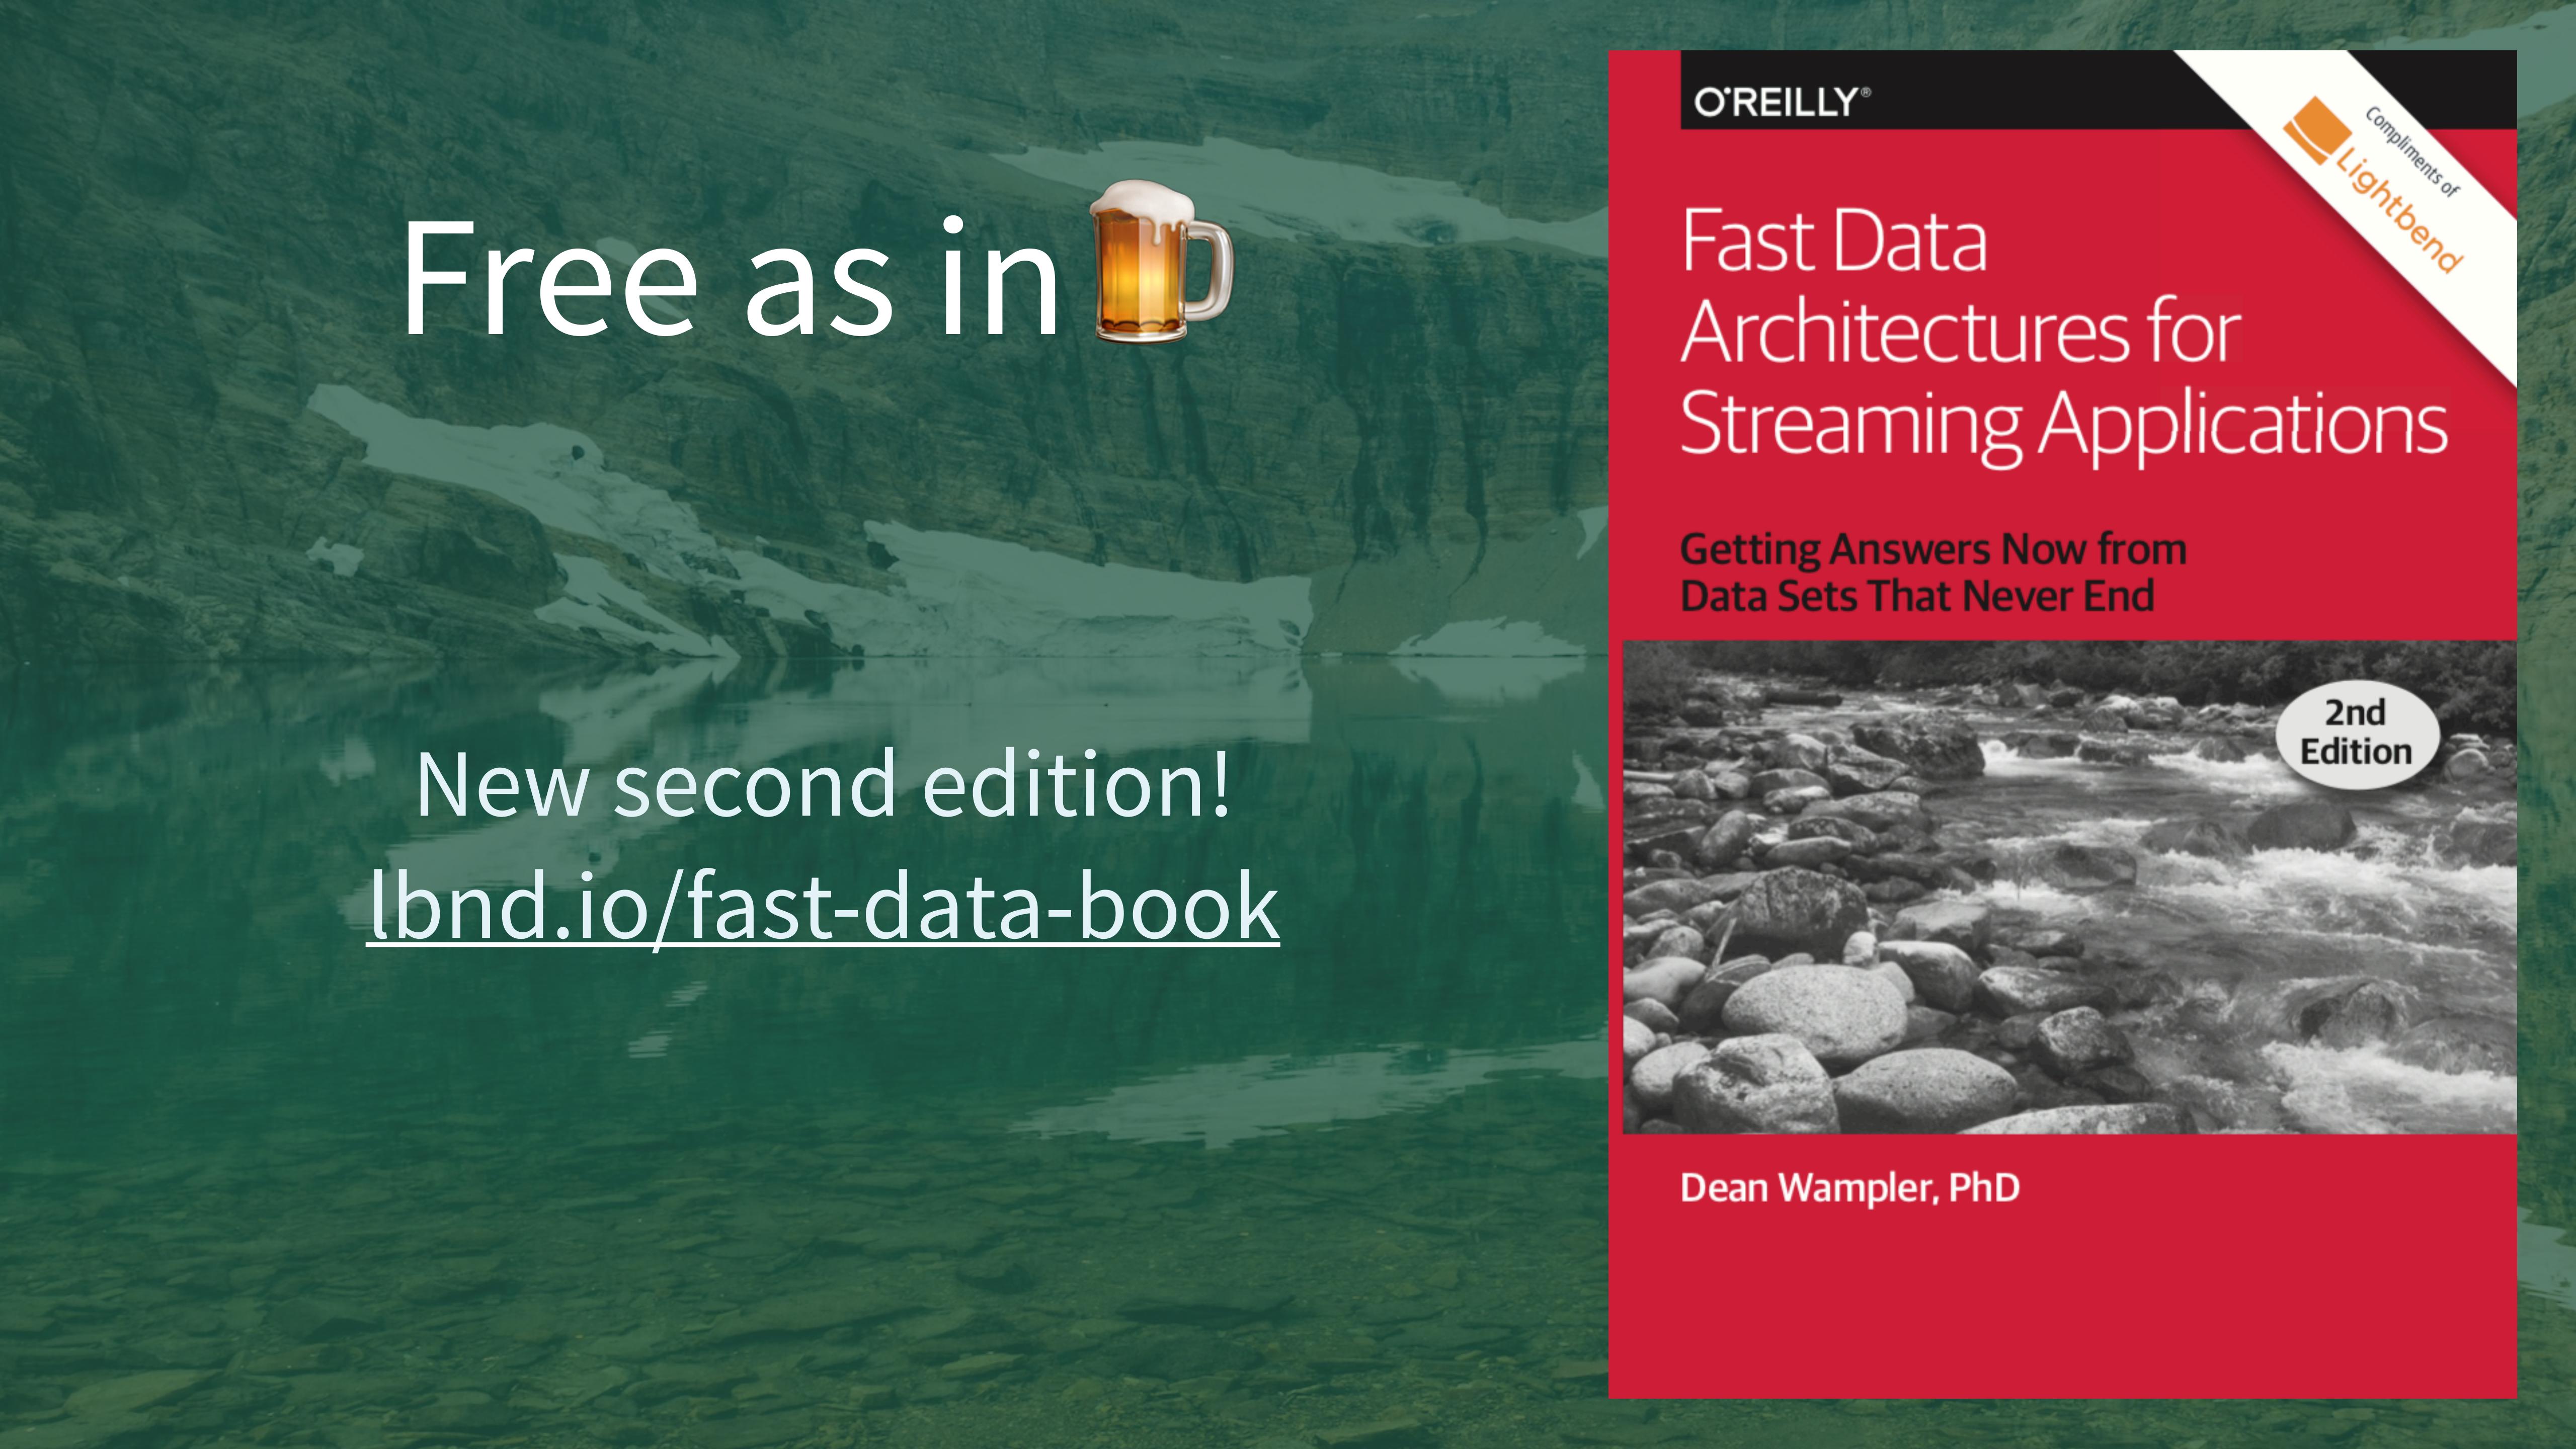
akka Streams

Streams

Event-driven μ-services

“Record-centric” μ-services





Free as in 🍺

New second edition!
lbnd.io/fast-data-book

O'REILLY®

Compliments of
Lightbend

Fast Data Architectures for Streaming Applications

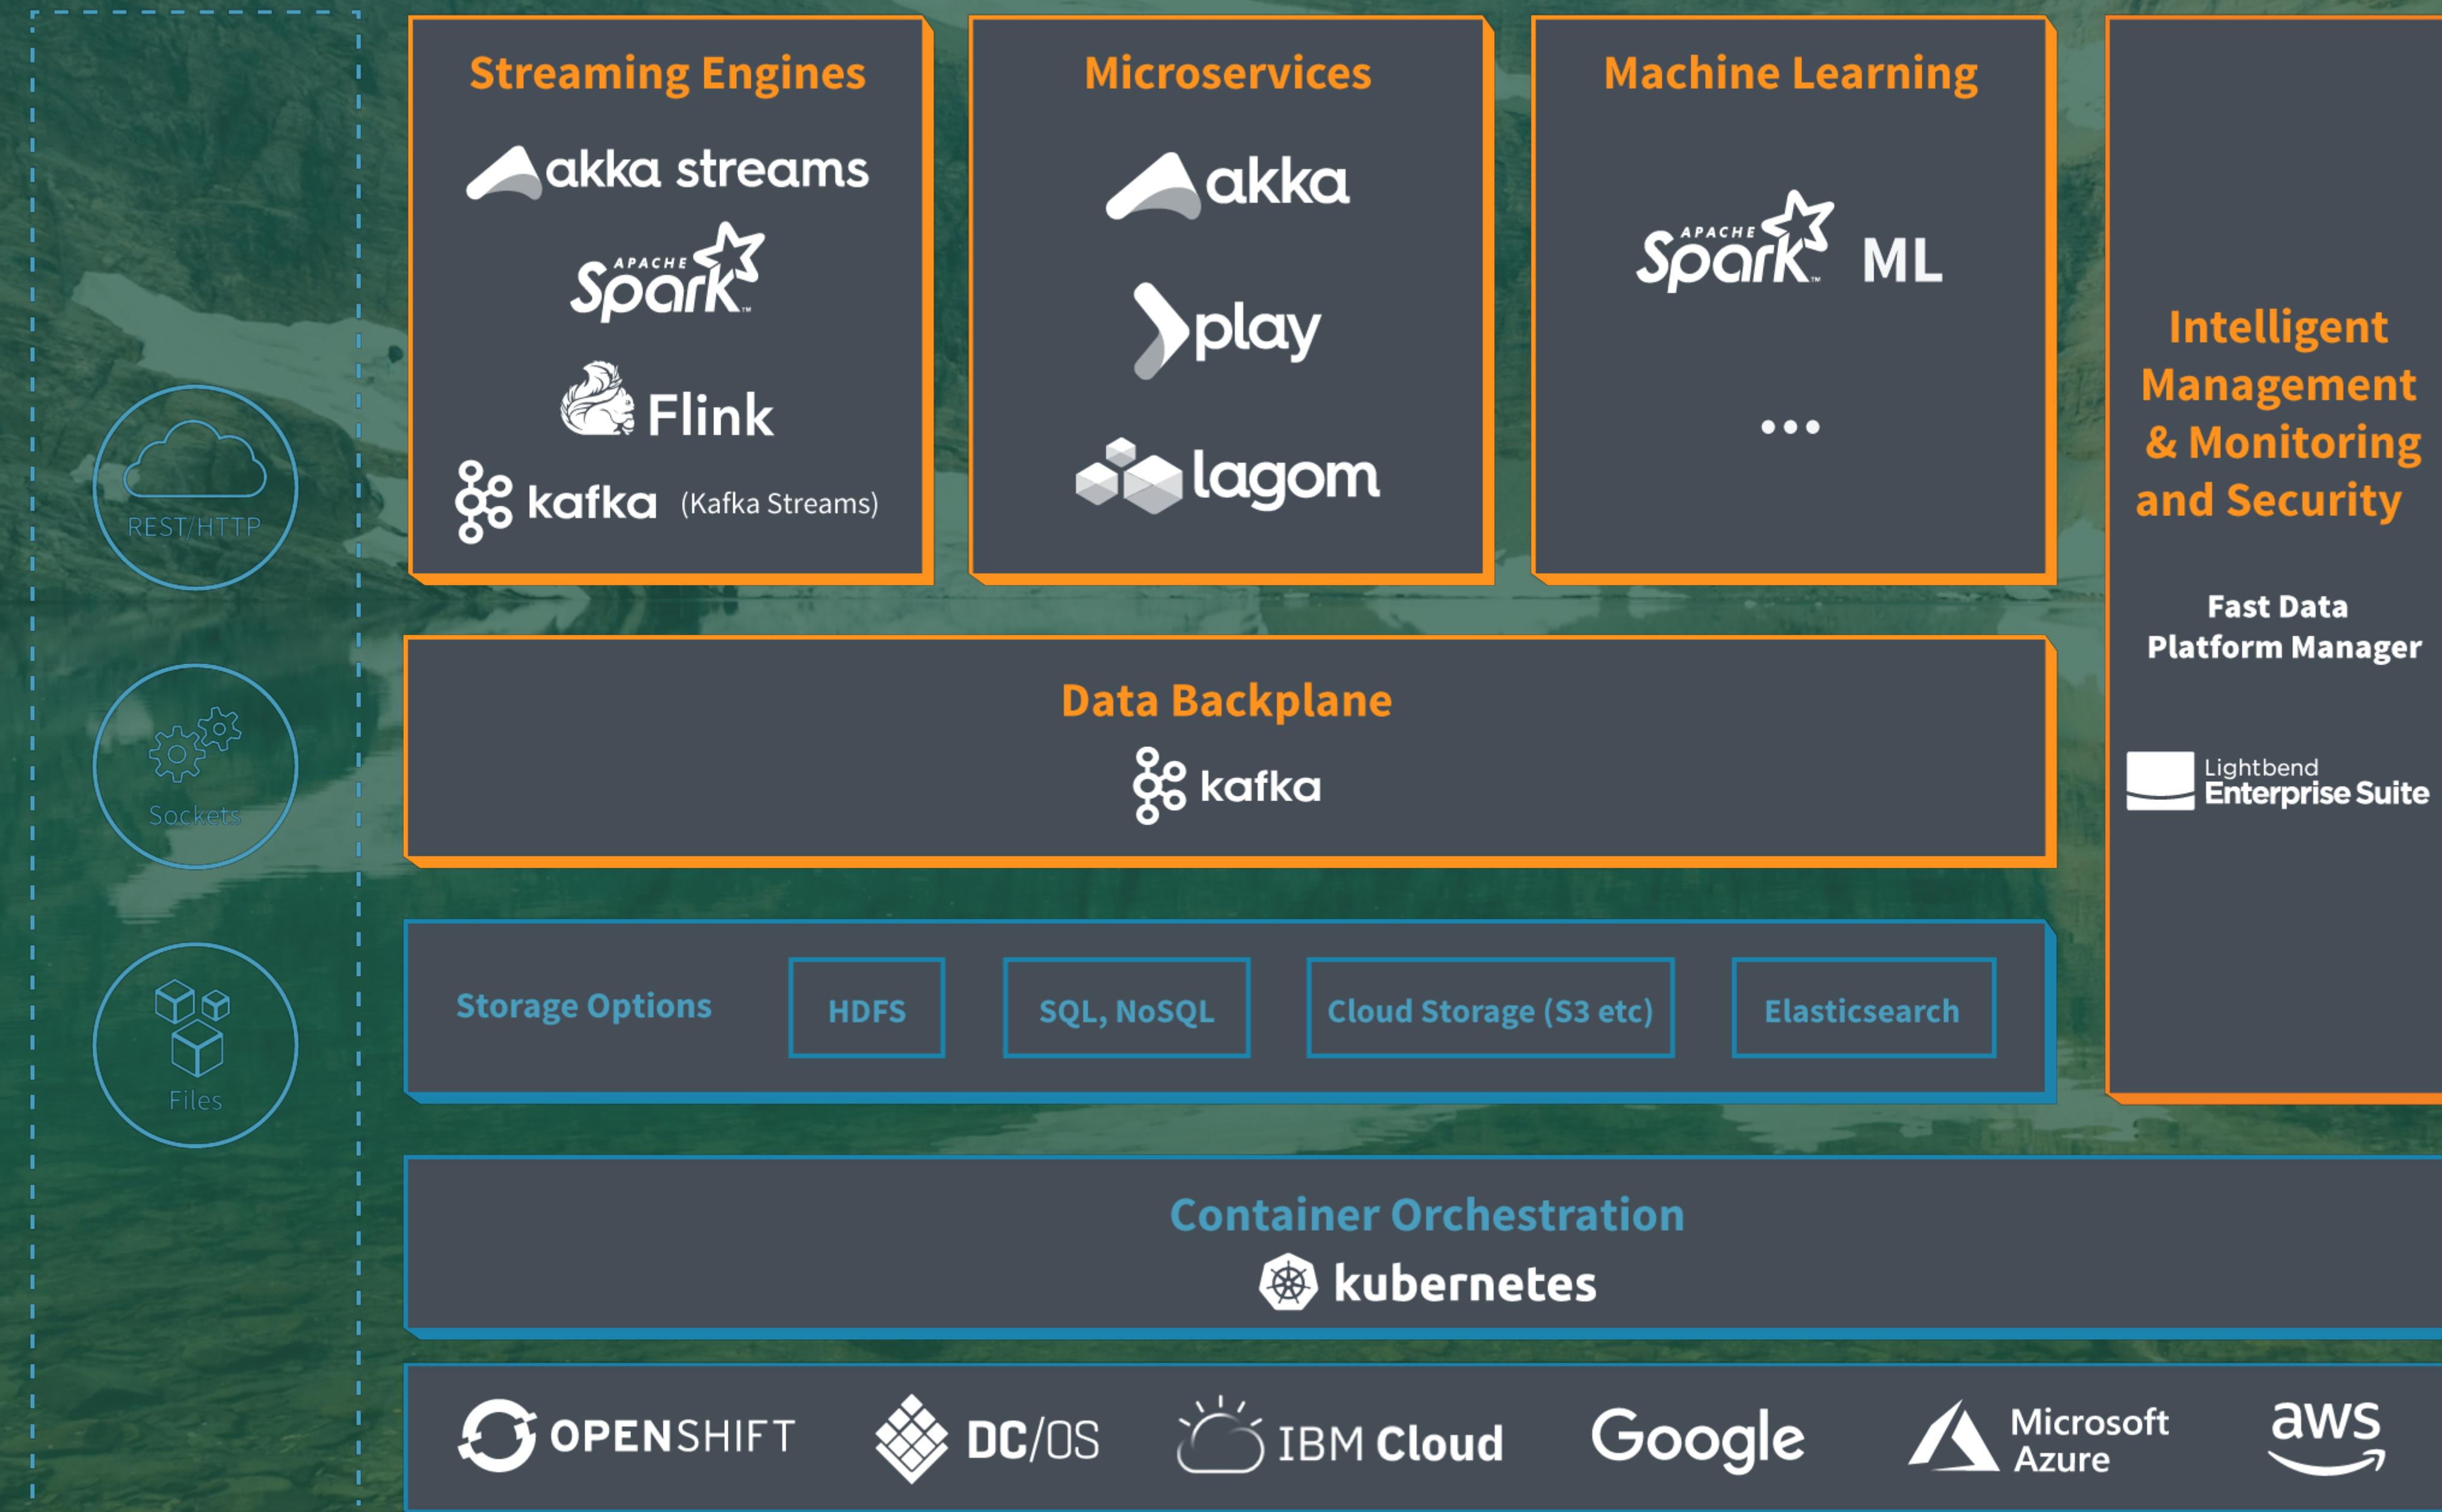
Getting Answers Now from
Data Sets That Never End

2nd
Edition



Dean Wampler, PhD

lightbend.com/fast-data-platform



Questions?

@deanwampler
lightbend.com/fast-data-platform
polyglotprogramming.com/talks