

The background of the slide features a stunning landscape of a lake and mountains at sunset. The sky is clear and blue, transitioning to a warm orange and yellow glow on the horizon. The mountains are rugged and rocky, with exposed layers of sedimentary rock. The lake in the foreground is calm, reflecting the colors of the sky and the surrounding mountains. The overall atmosphere is serene and majestic.

# Streaming Data Microservices

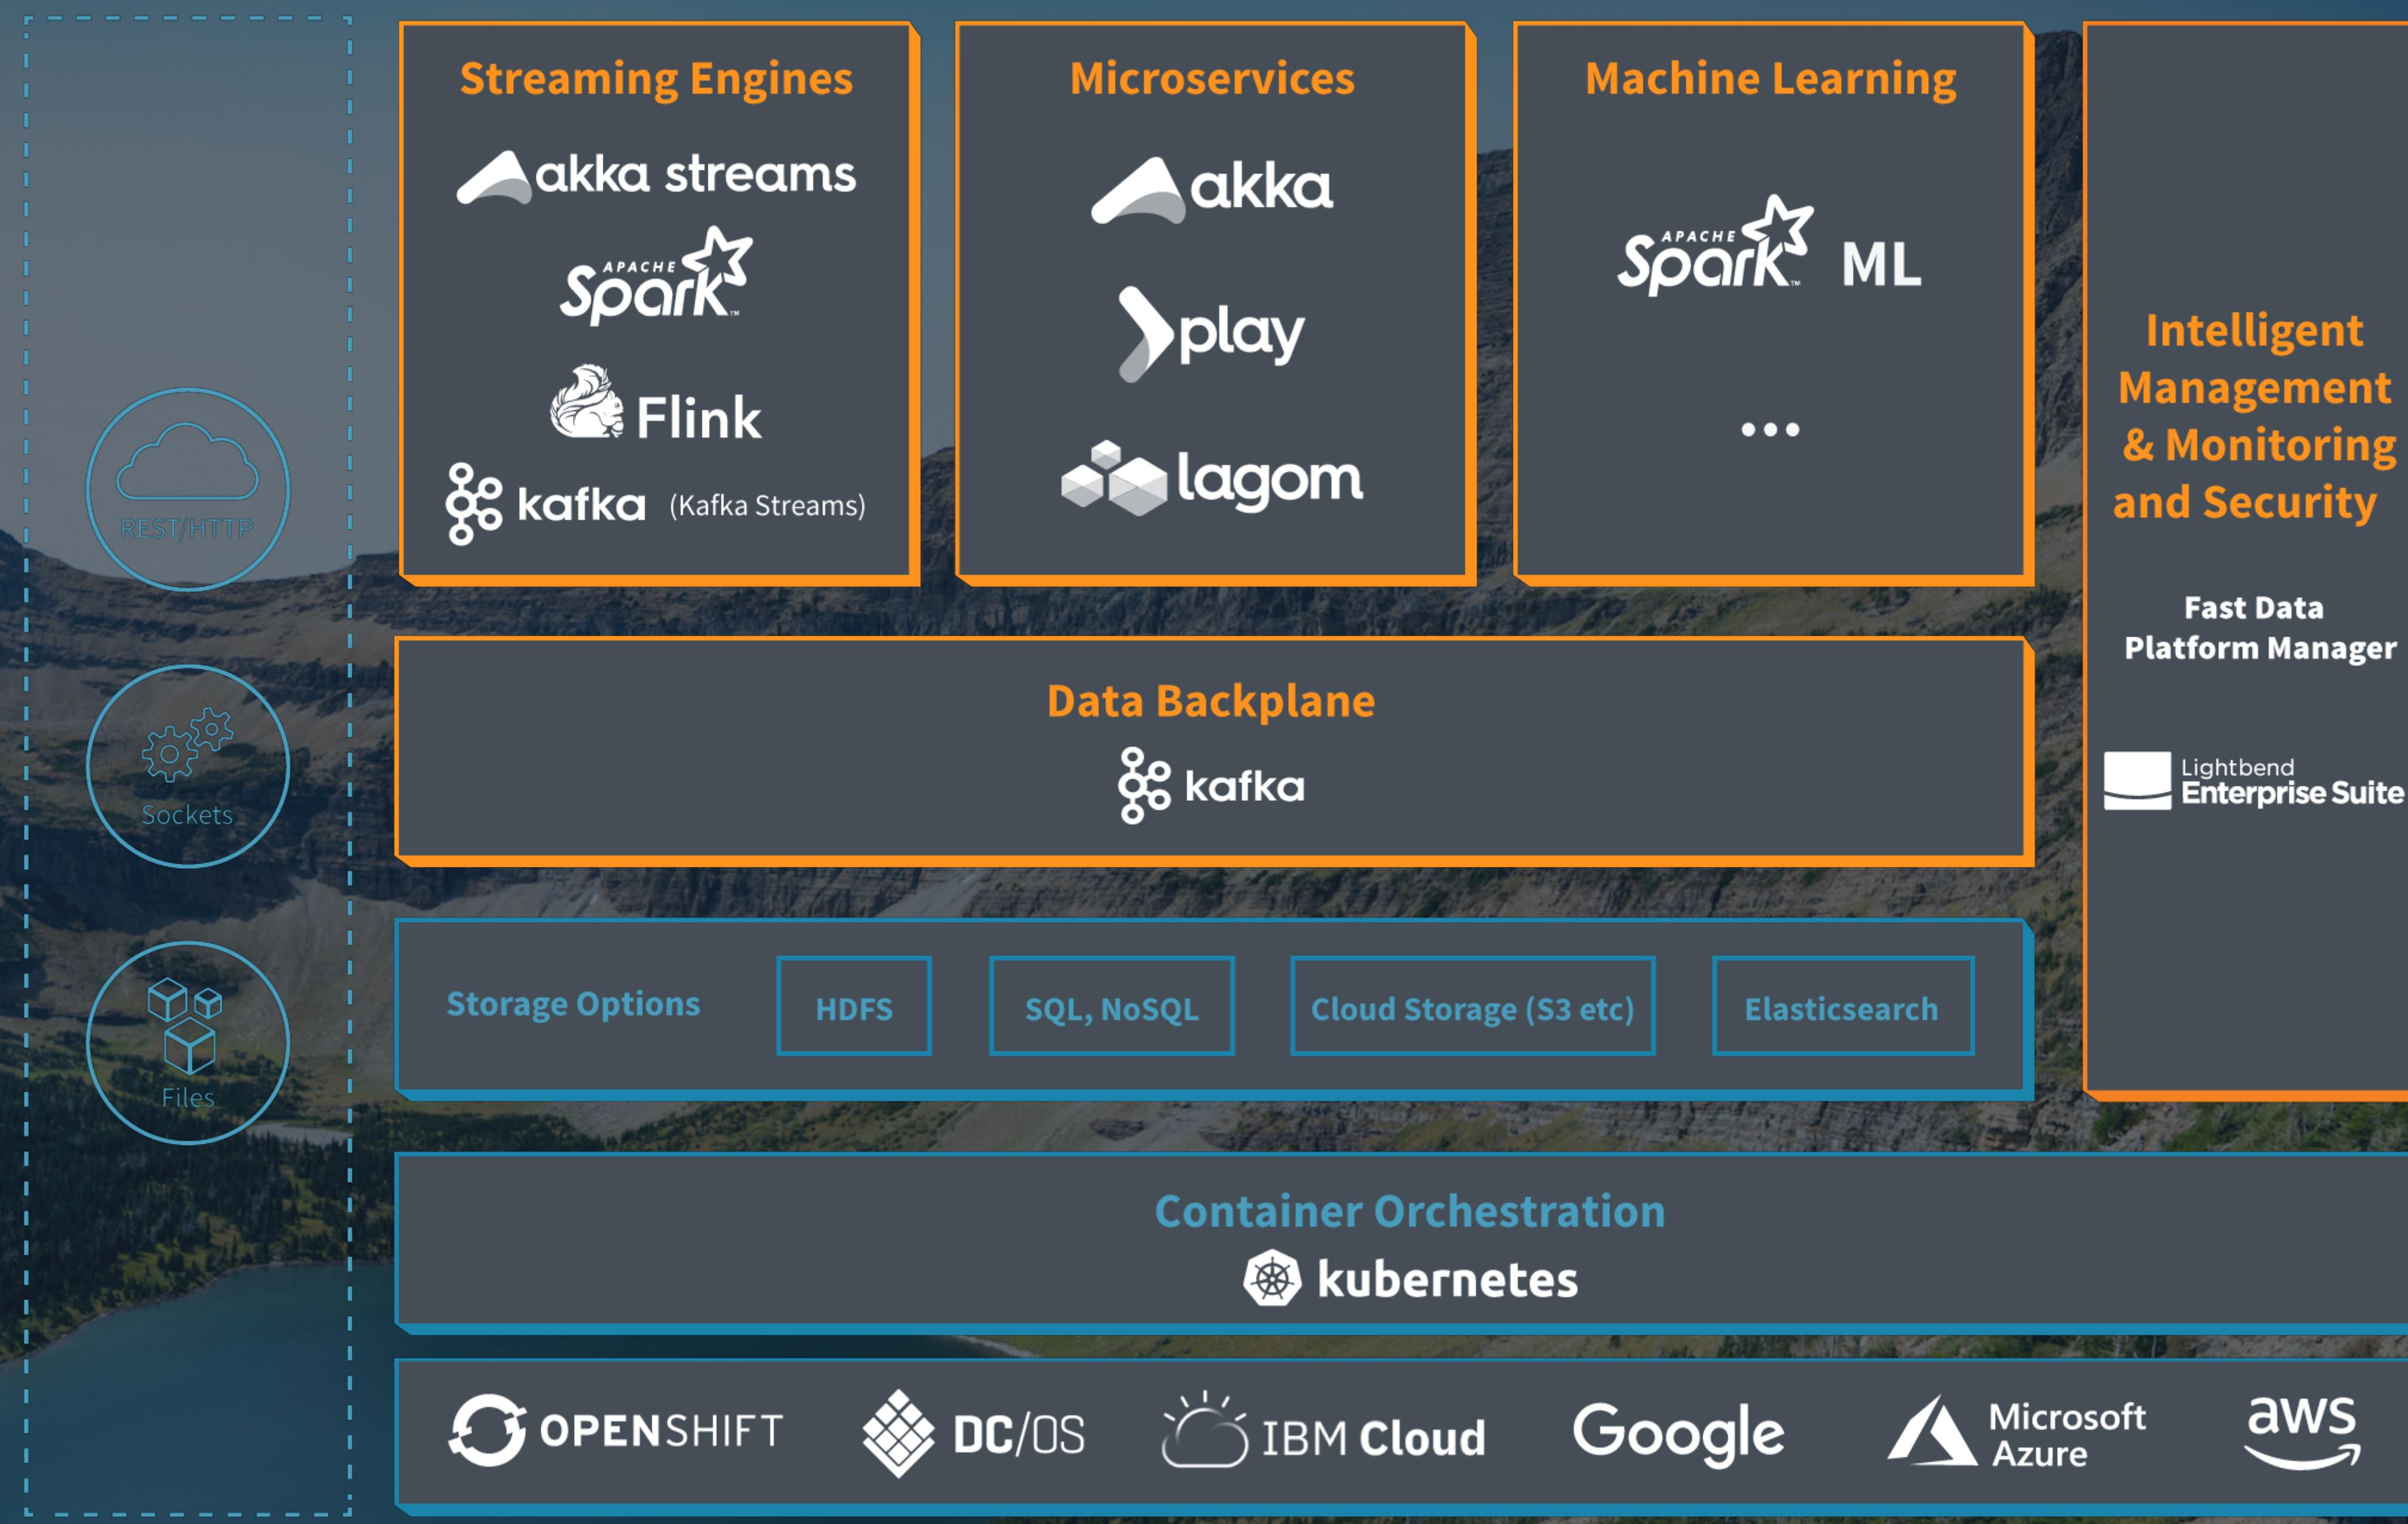
## With Akka Streams and Kafka Streams

Dean Wampler, Ph.D.  
[dean@lightbend.com](mailto:dean@lightbend.com)  
[@deanwampler](https://twitter.com/deanwampler)

# Who am I?



@deanwampler



I lead the Lightbend Fast Data Platform; unified streaming data & microservices

# [lightbend.com/fast-data-platform](http://lightbend.com/fast-data-platform)

## Streaming Engines

akka streams



kafka (Kafka Streams)



## Microservices



lagom

## Machine Learning



...

Intelligent  
Management  
& Monitoring  
and Security

Fast Data  
Platform Manager

Lightbend  
Enterprise Suite

## Data Backplane



## Storage Options

HDFS

SQL, NoSQL

Cloud Storage (S3 etc)

Elasticsearch



## Container Orchestration





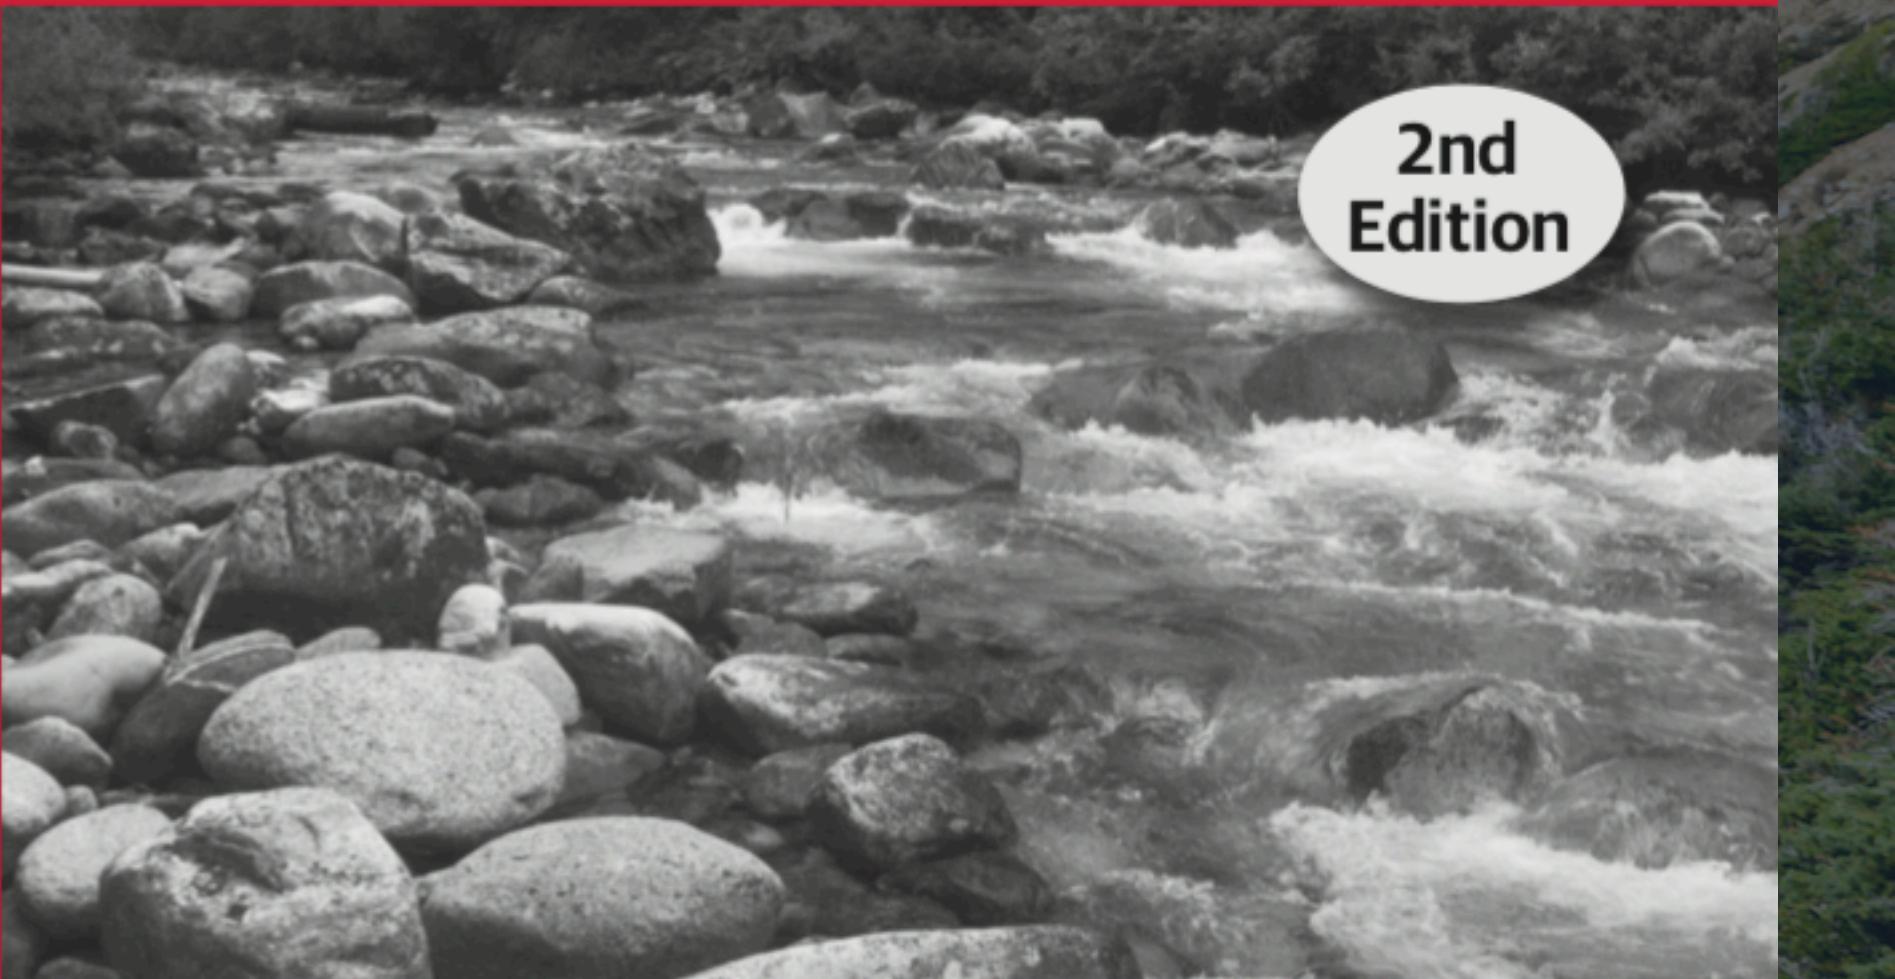
O'REILLY®

Free as in 

New second edition!  
[lbnd.io/fast-data-book](https://lbnd.io/fast-data-book)

# Fast Data Architectures for Streaming Applications

Getting Answers Now from  
Data Sets That Never End



2nd  
Edition

Dean Wampler, PhD

TL;DR



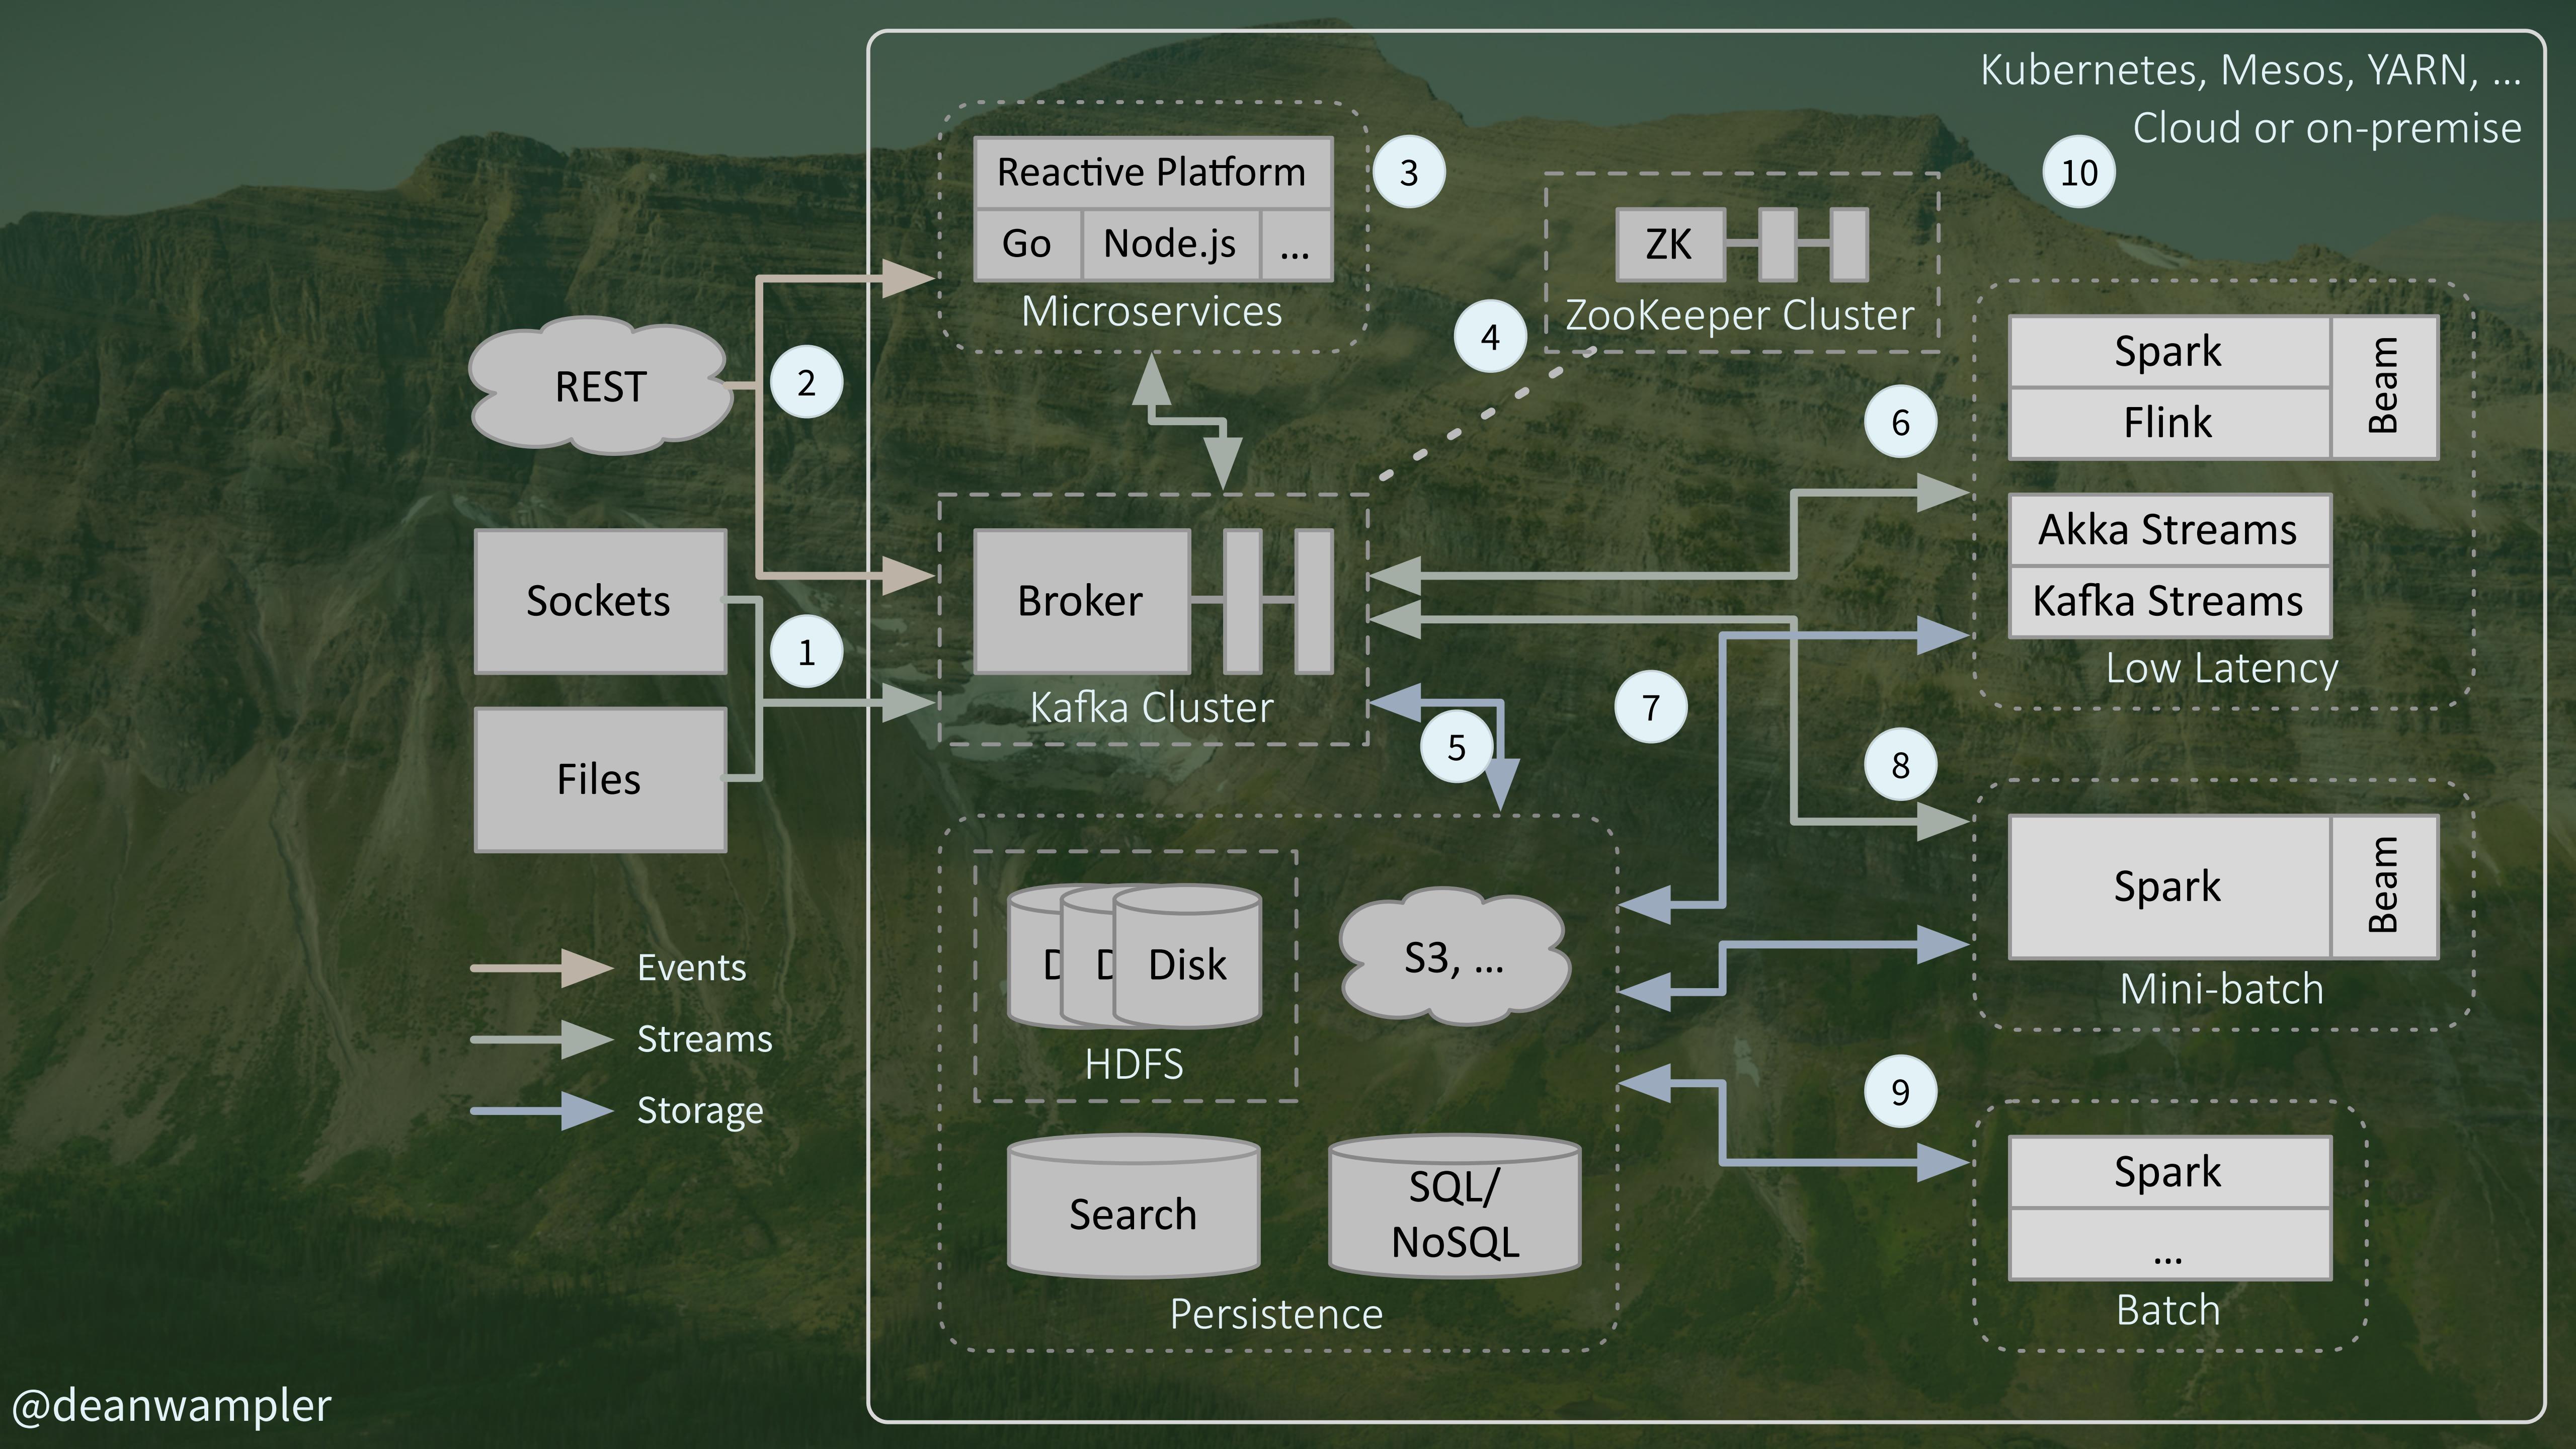
# TL;DR

For streaming data:

- Use Kafka as the “backplane”
- Use Spark or Flink when you need massive scale...
- Use microservice *libraries*, like Akka Streams and Kafka Streams for everything else...



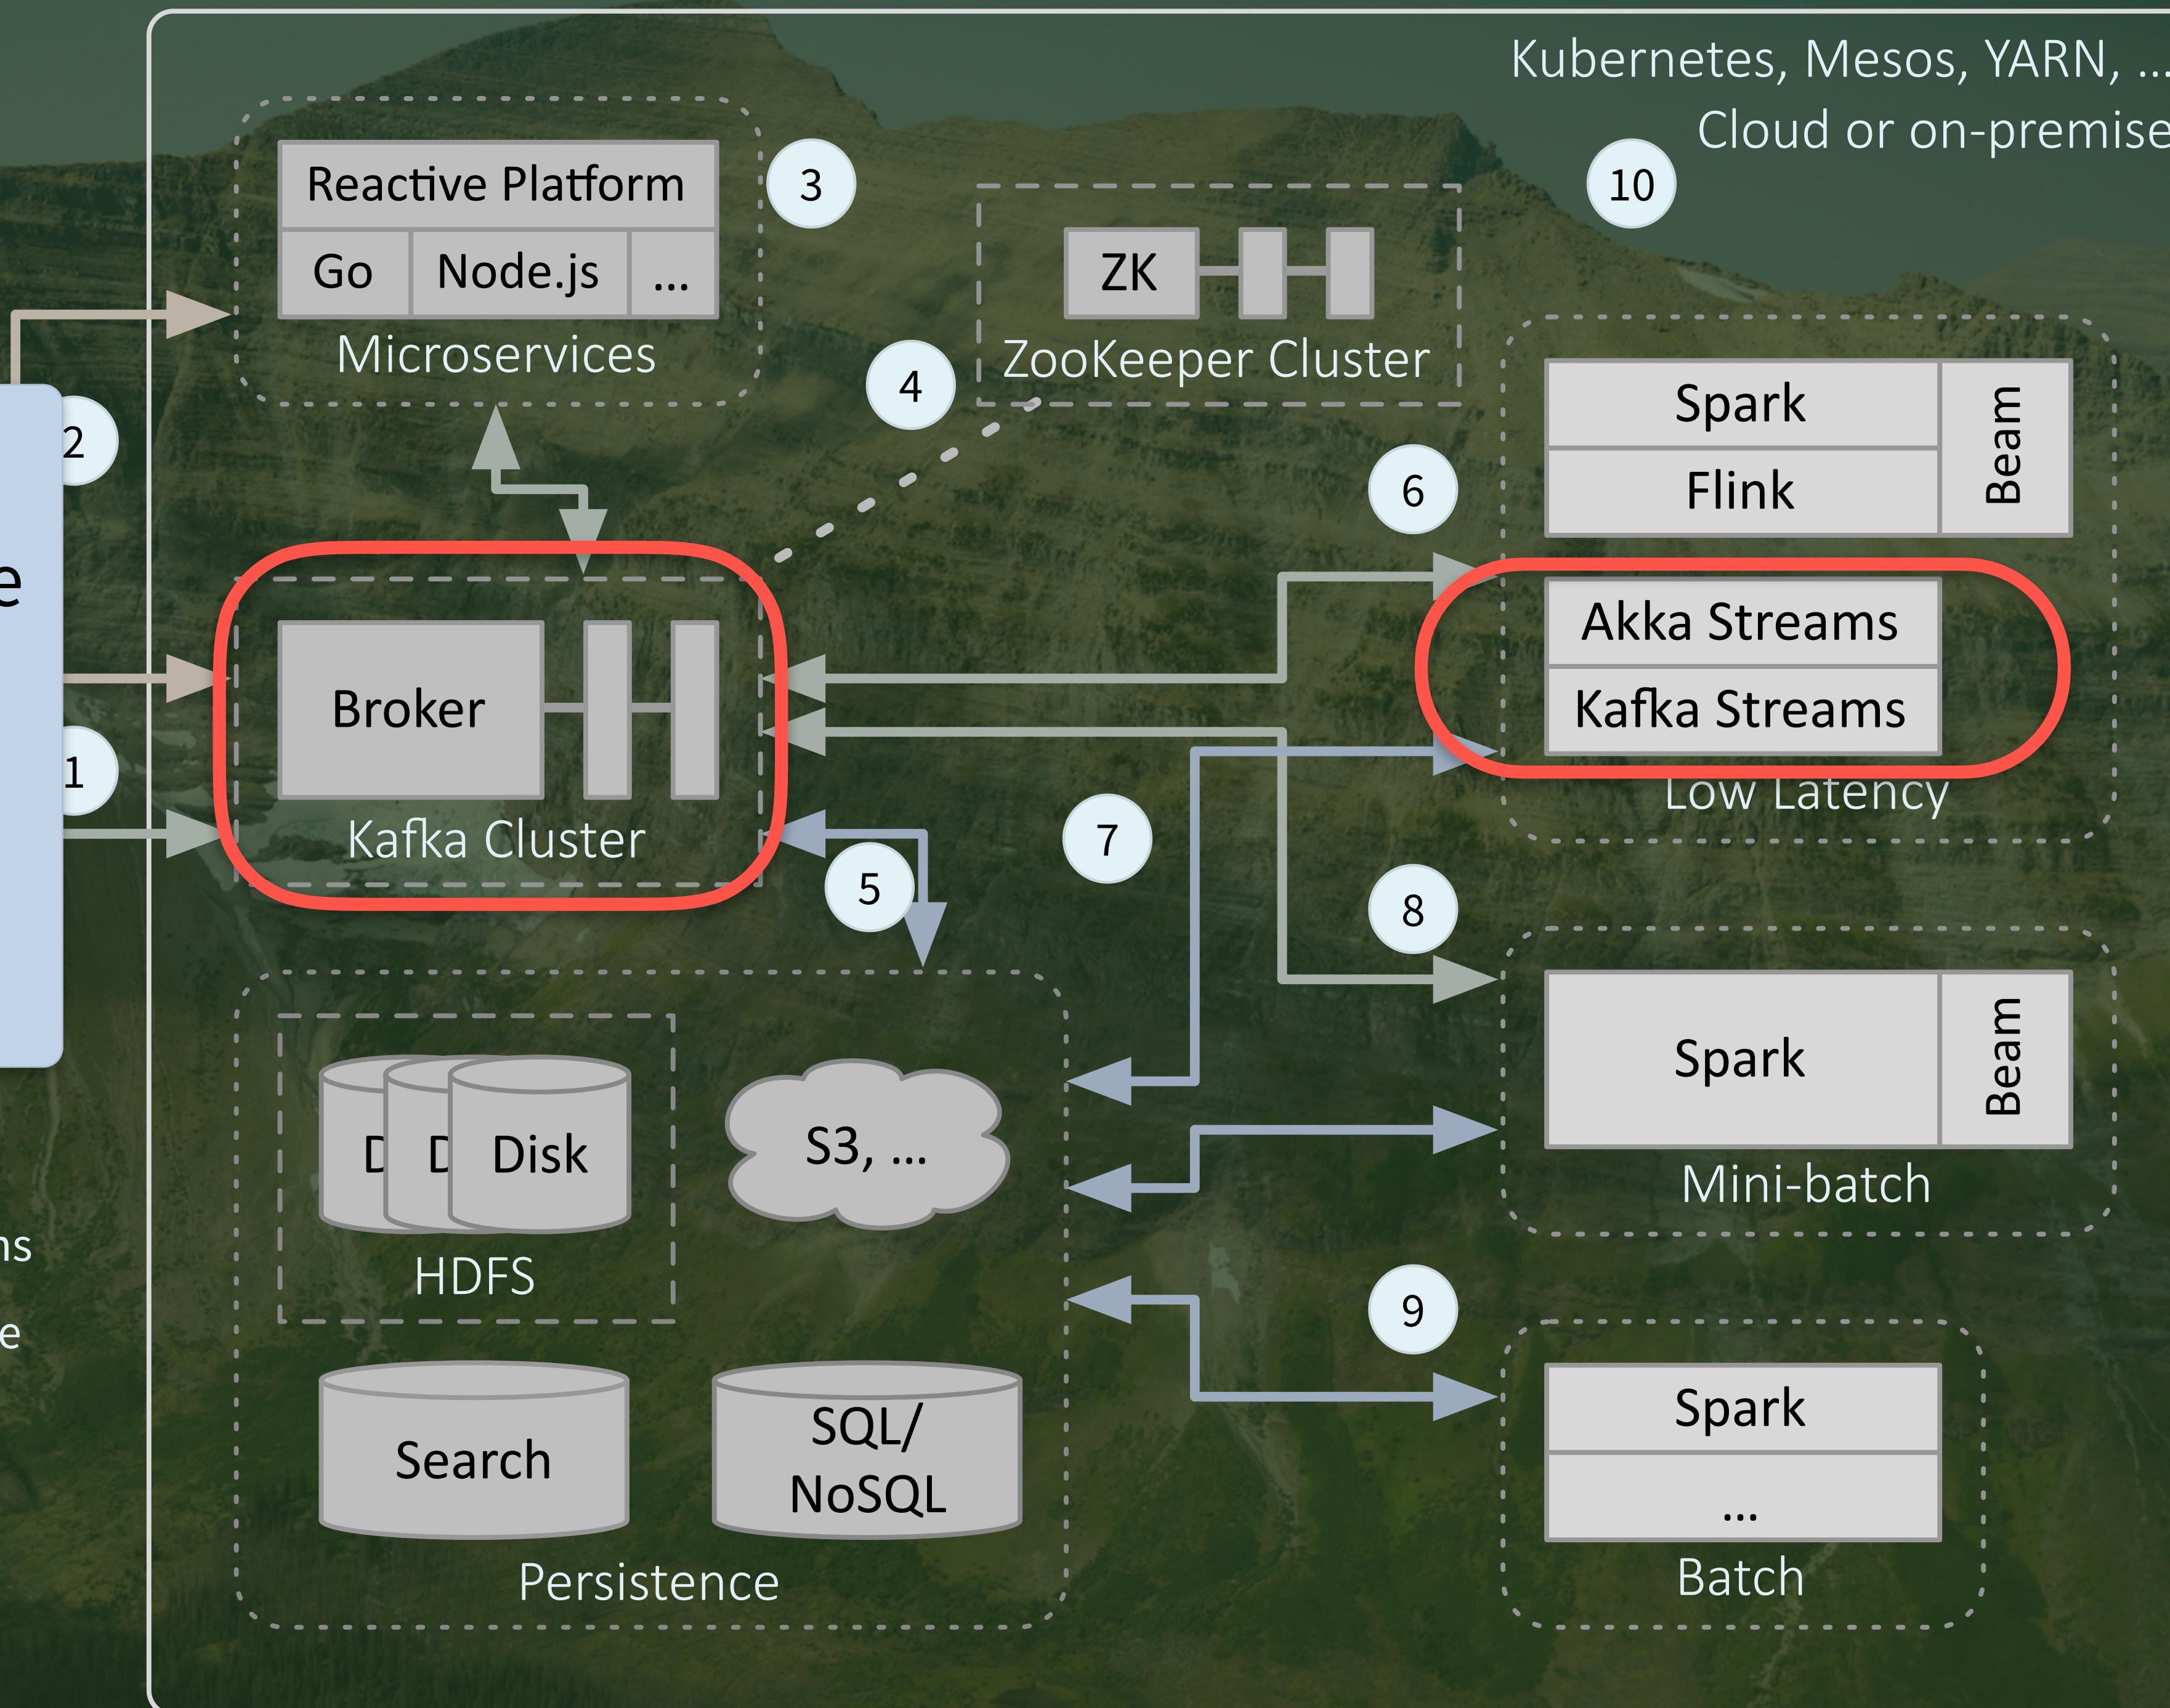
Streaming architectures (from the report)



This talk's focus:

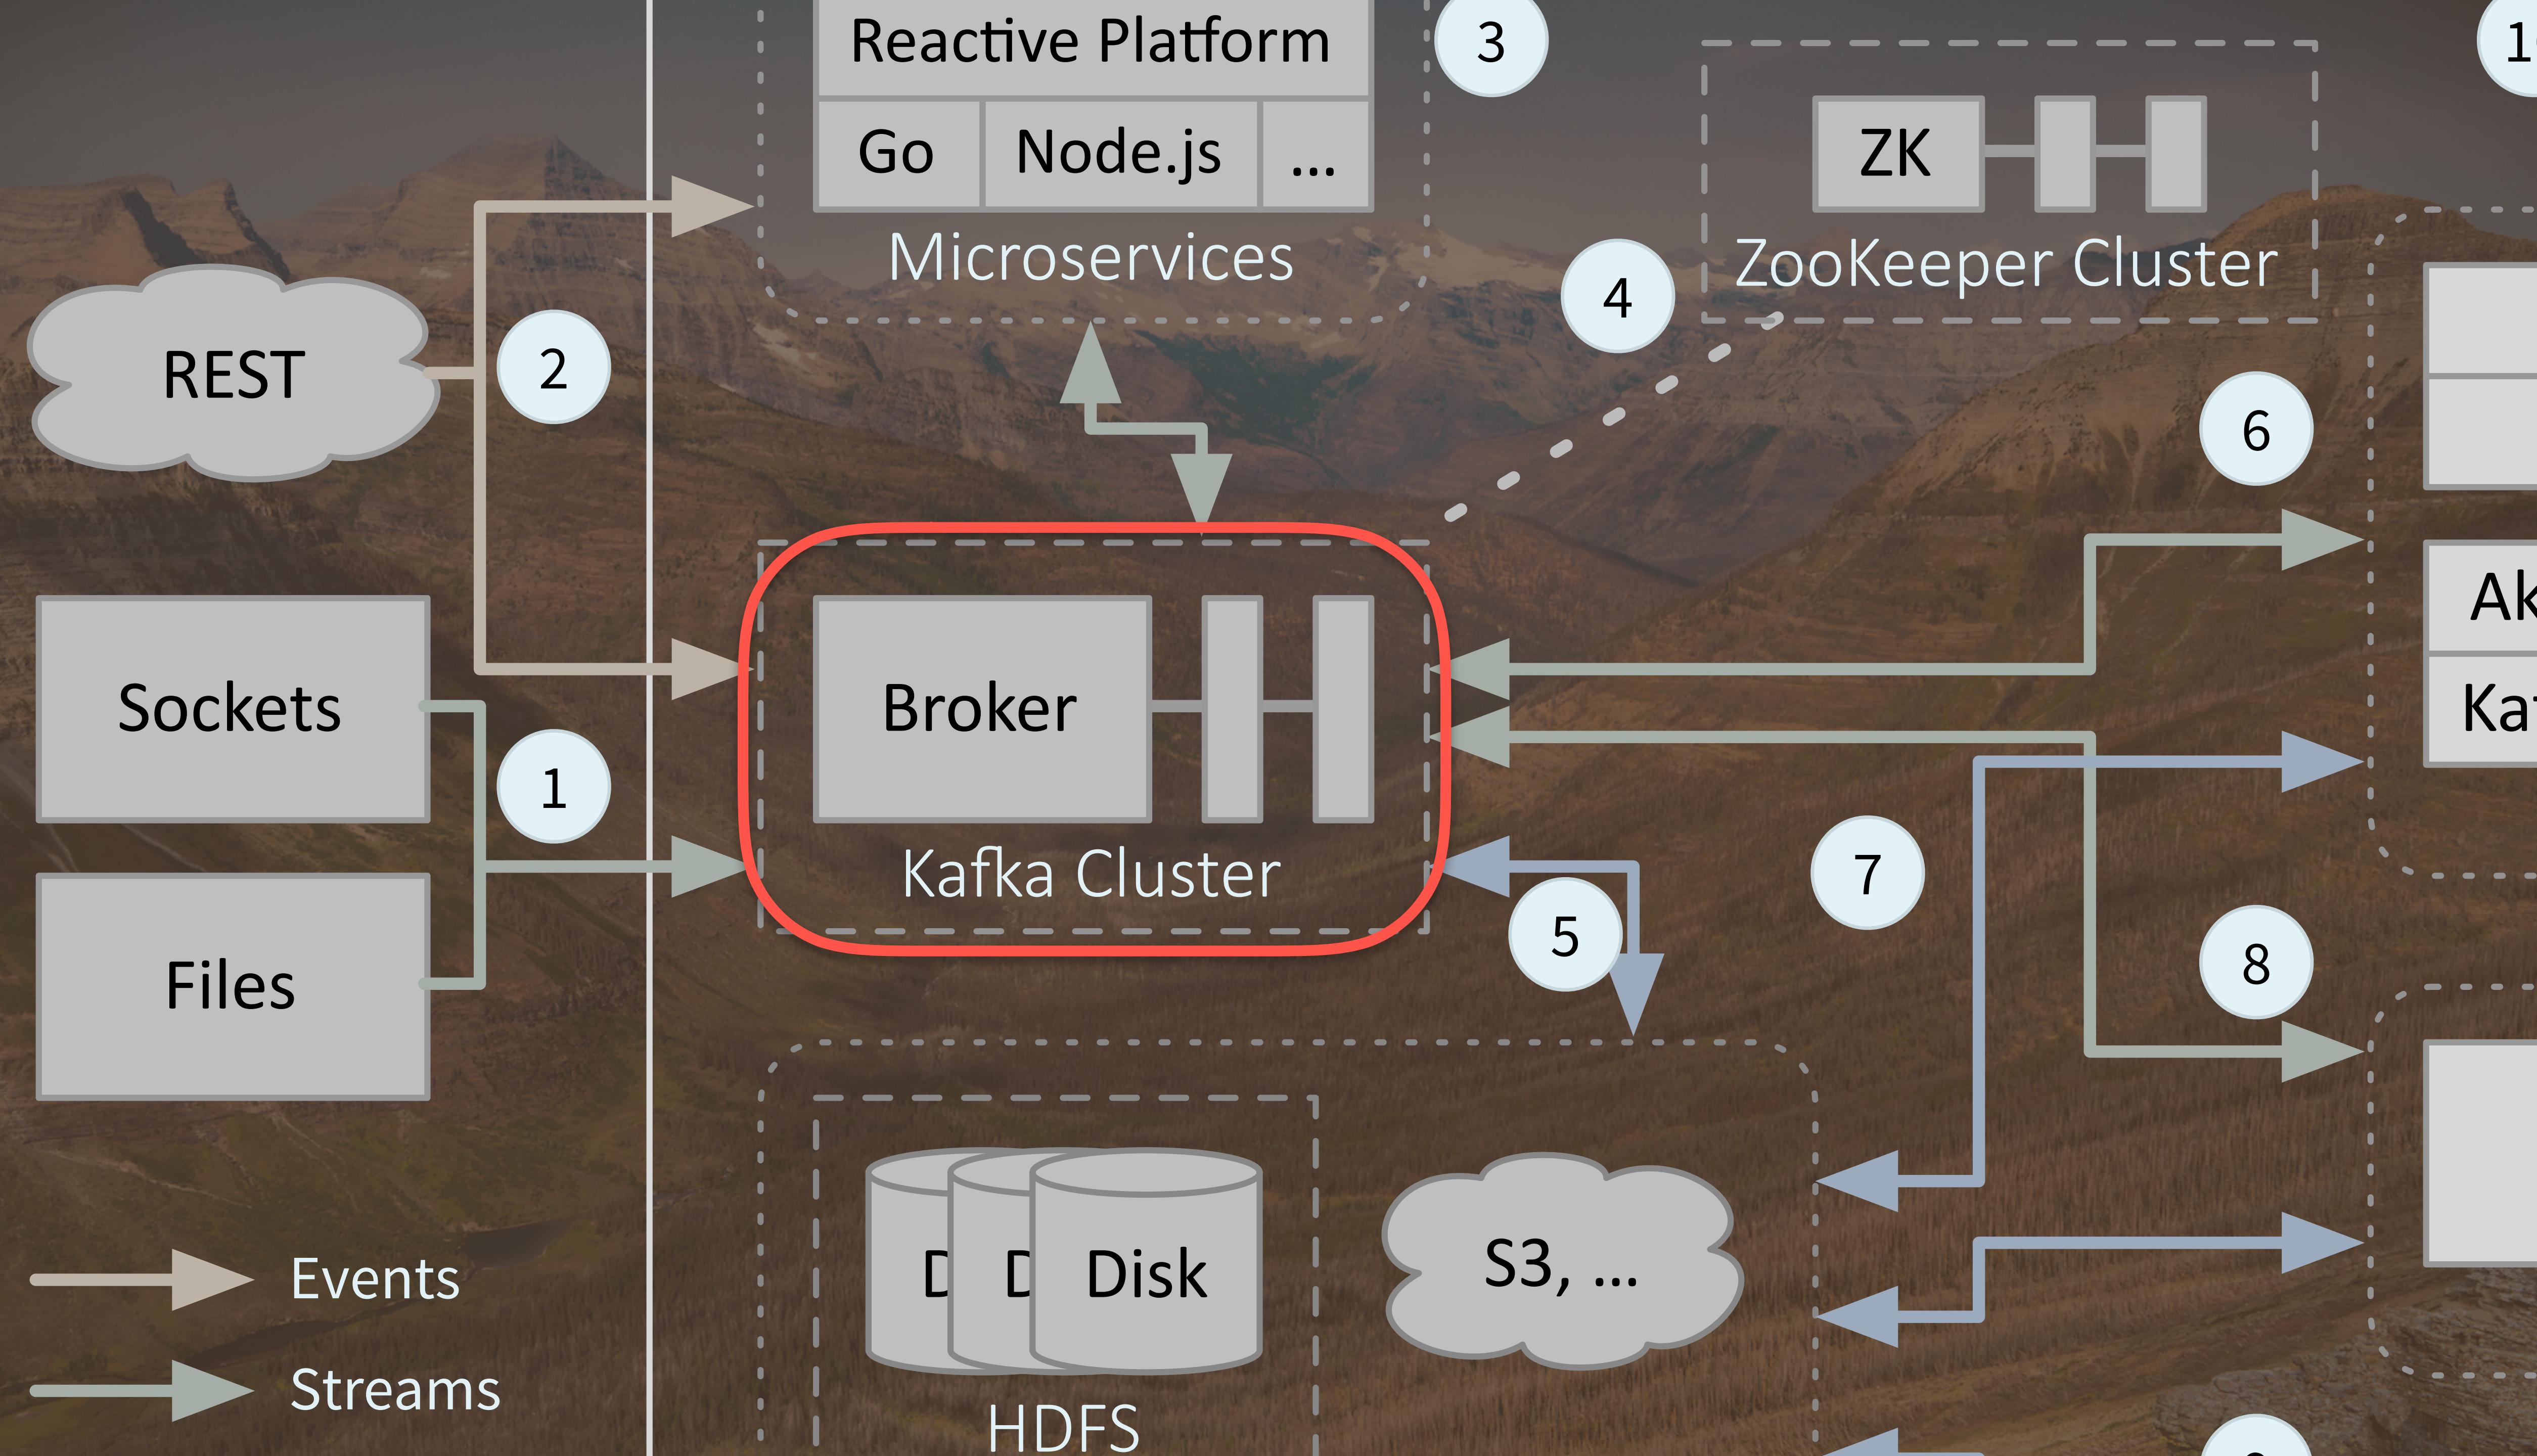
- Kafka - the data backplane
- Akka Streams and Kafka Streams - the engine for streaming data microservices

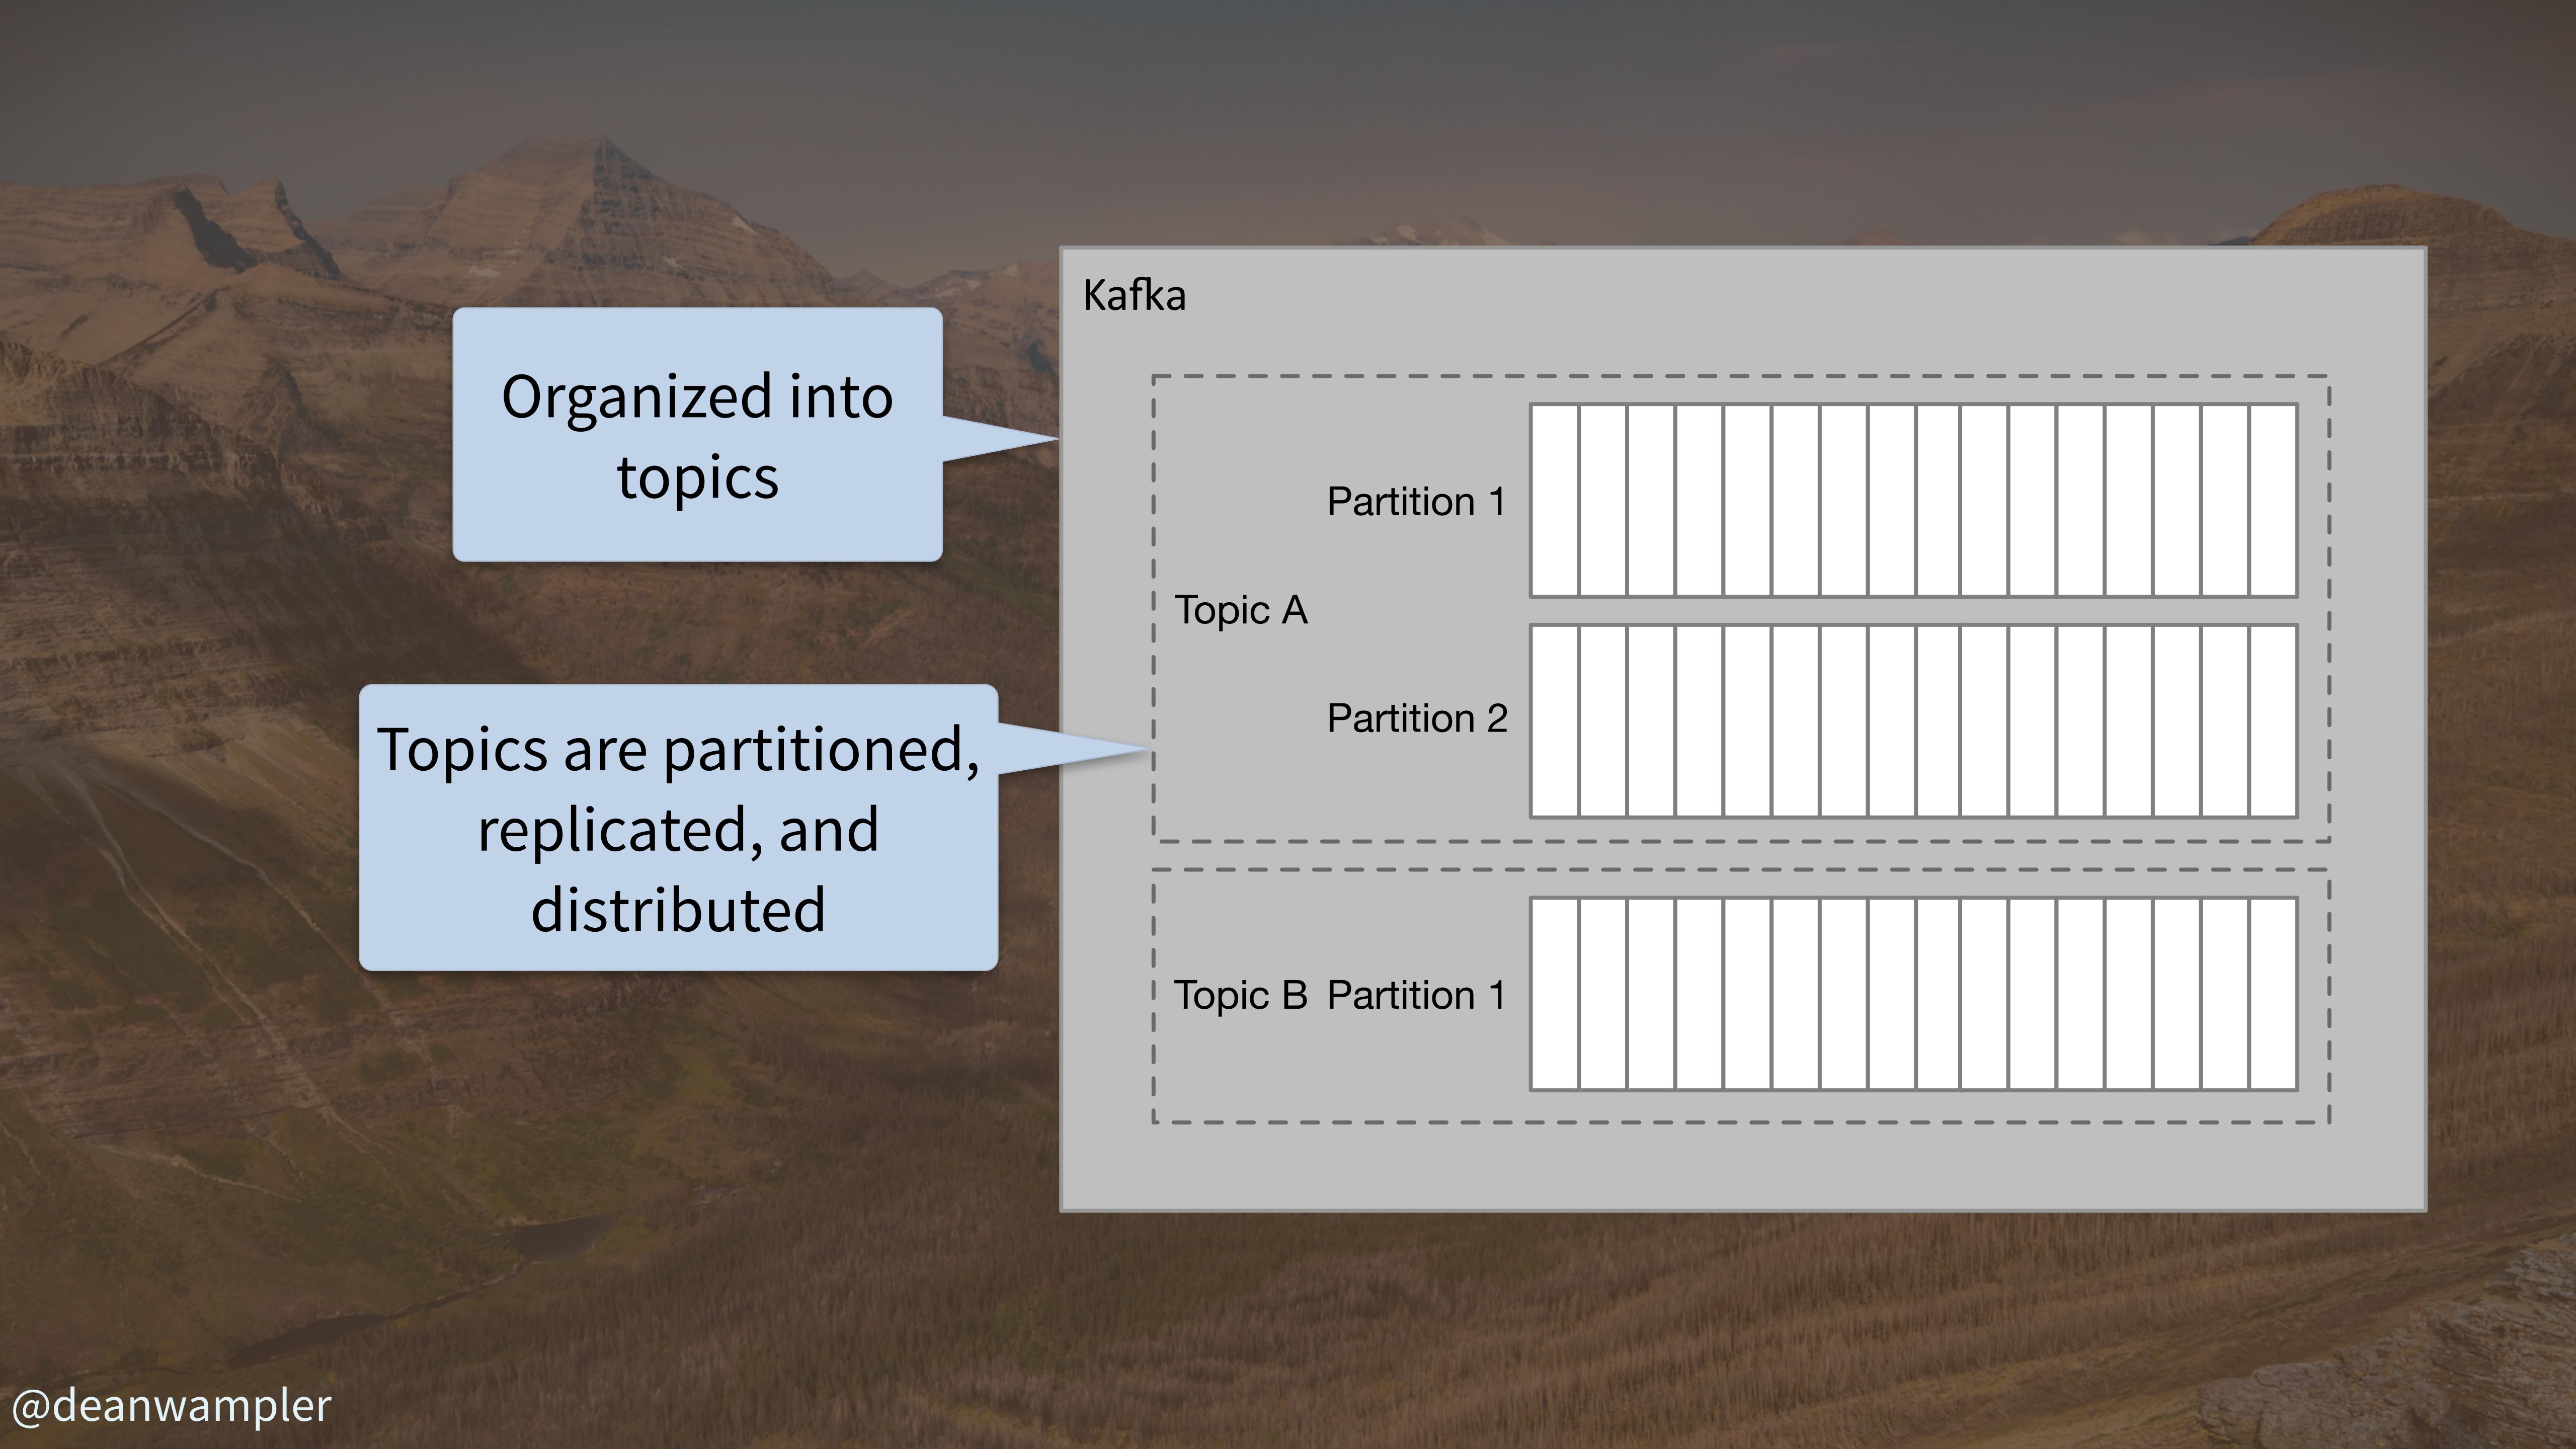
- Events
- Streams
- Storage



# Why Kafka?

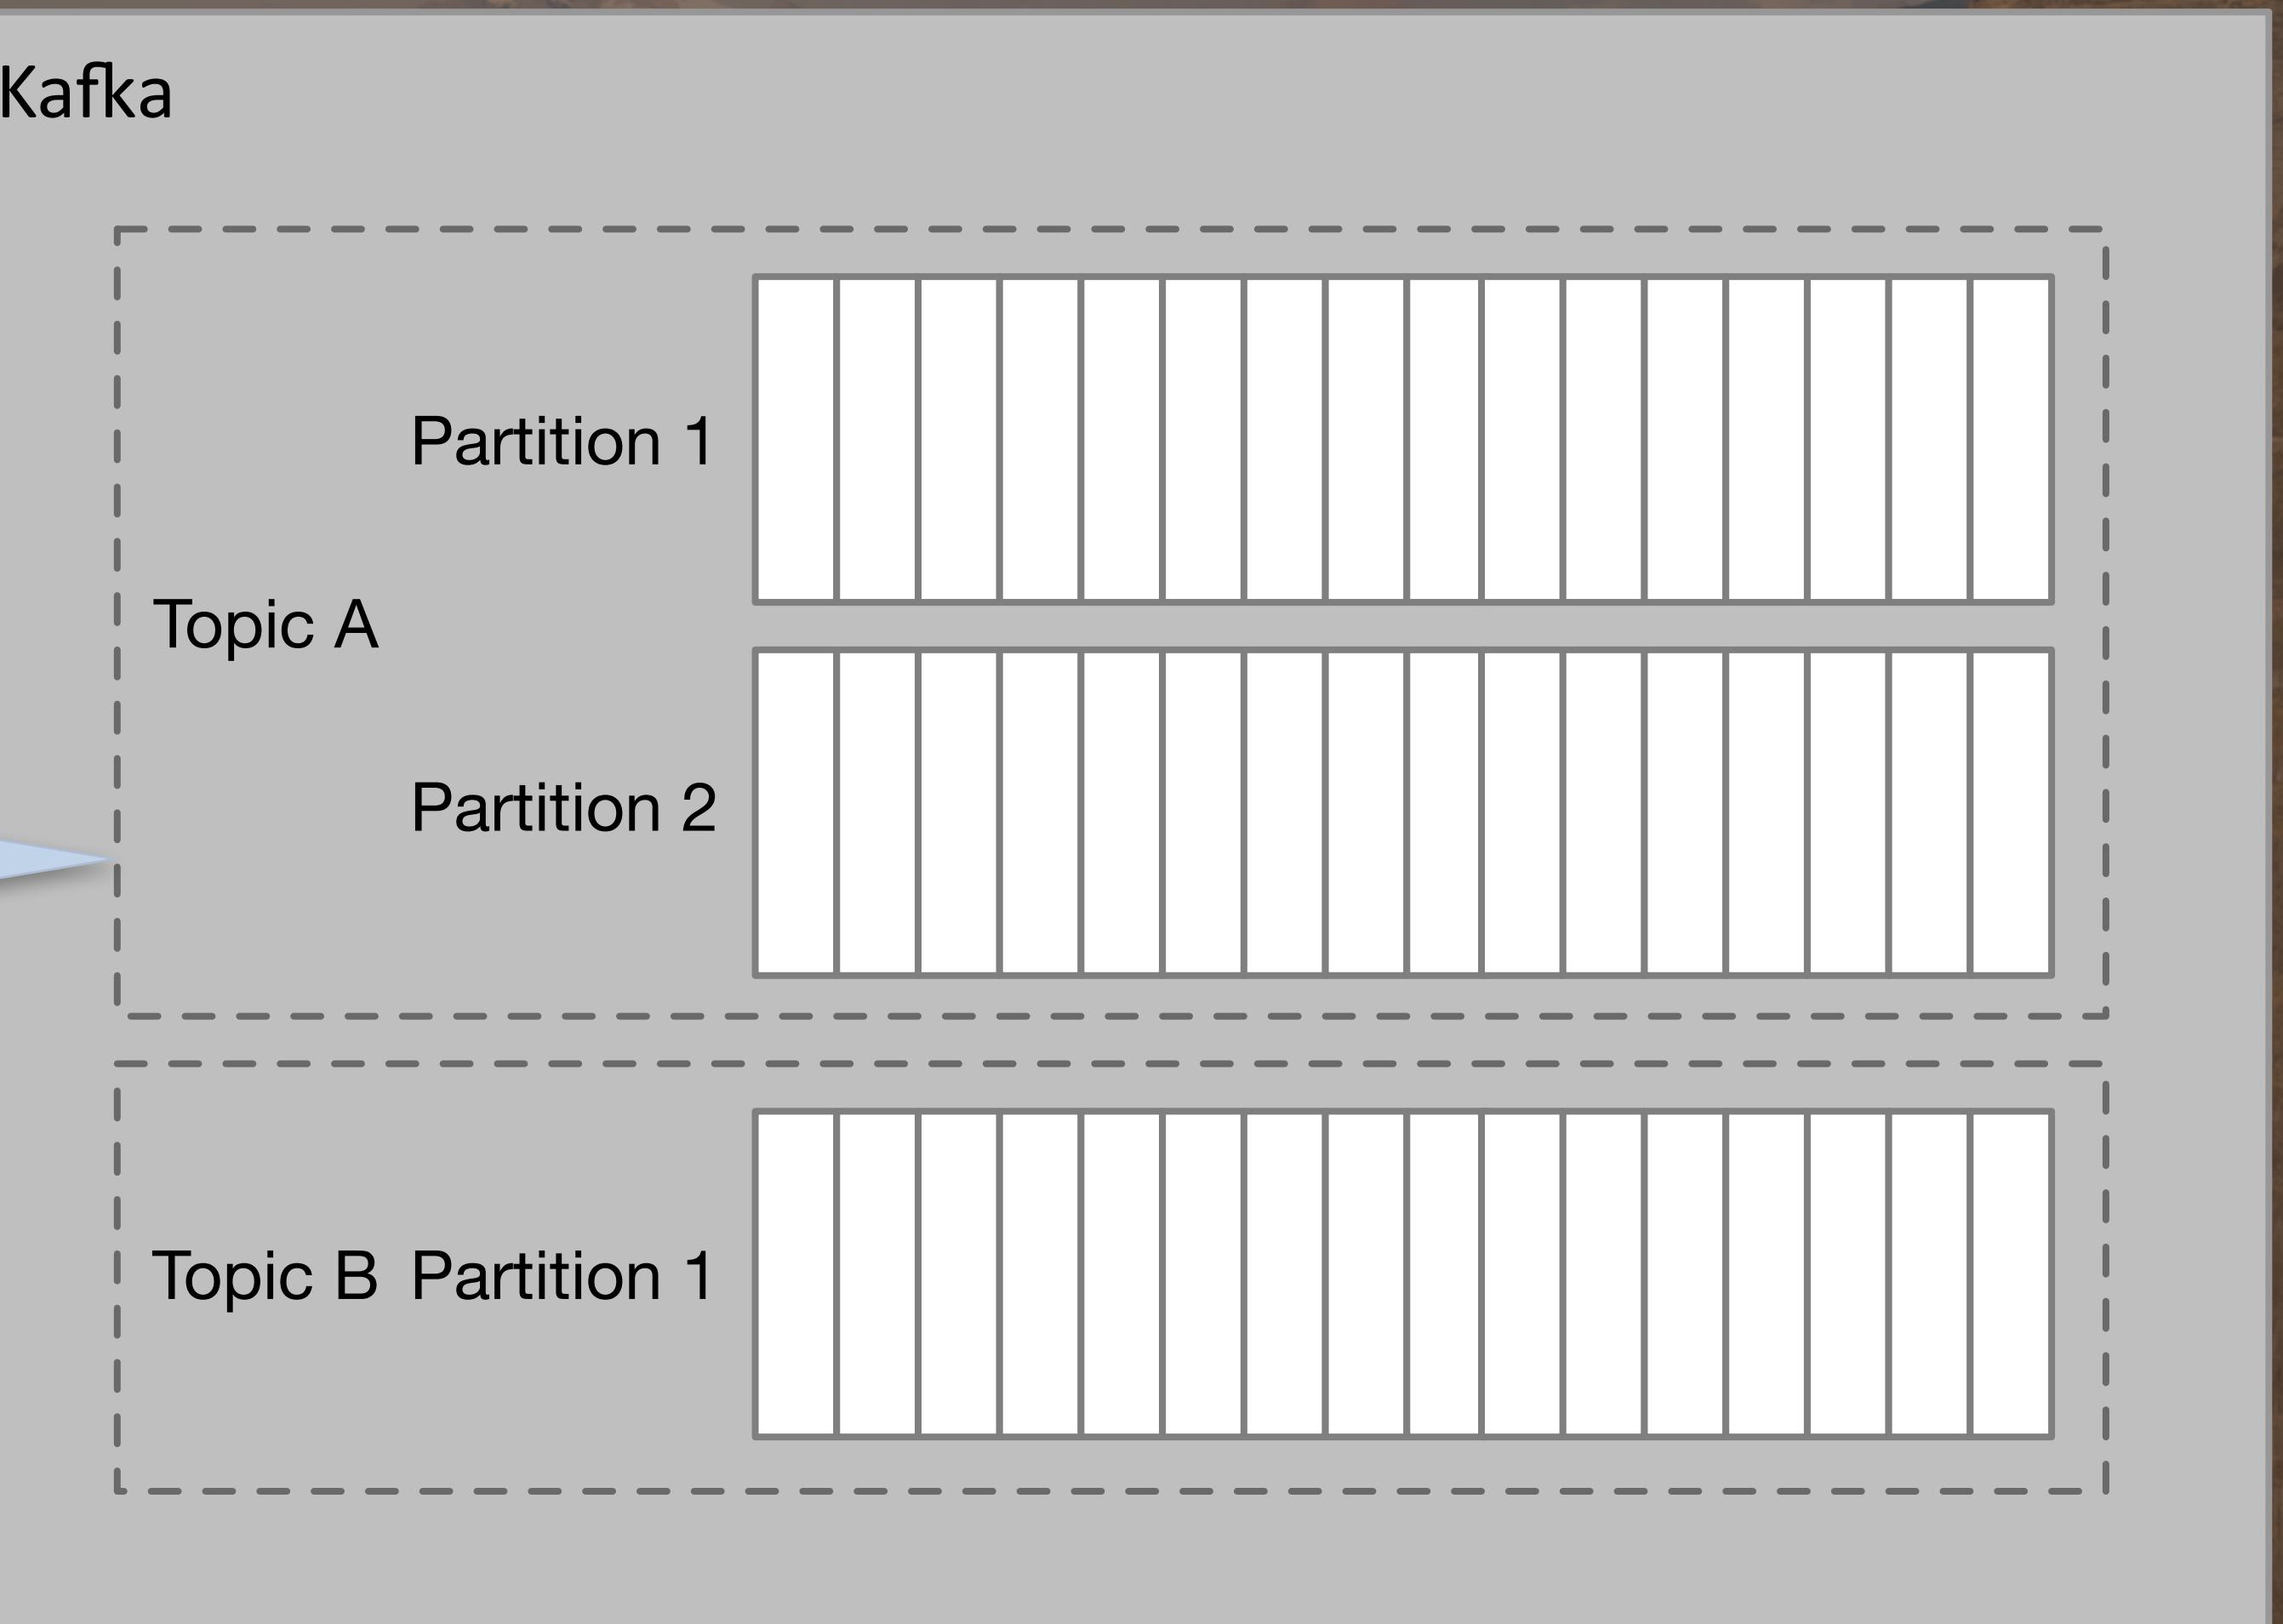




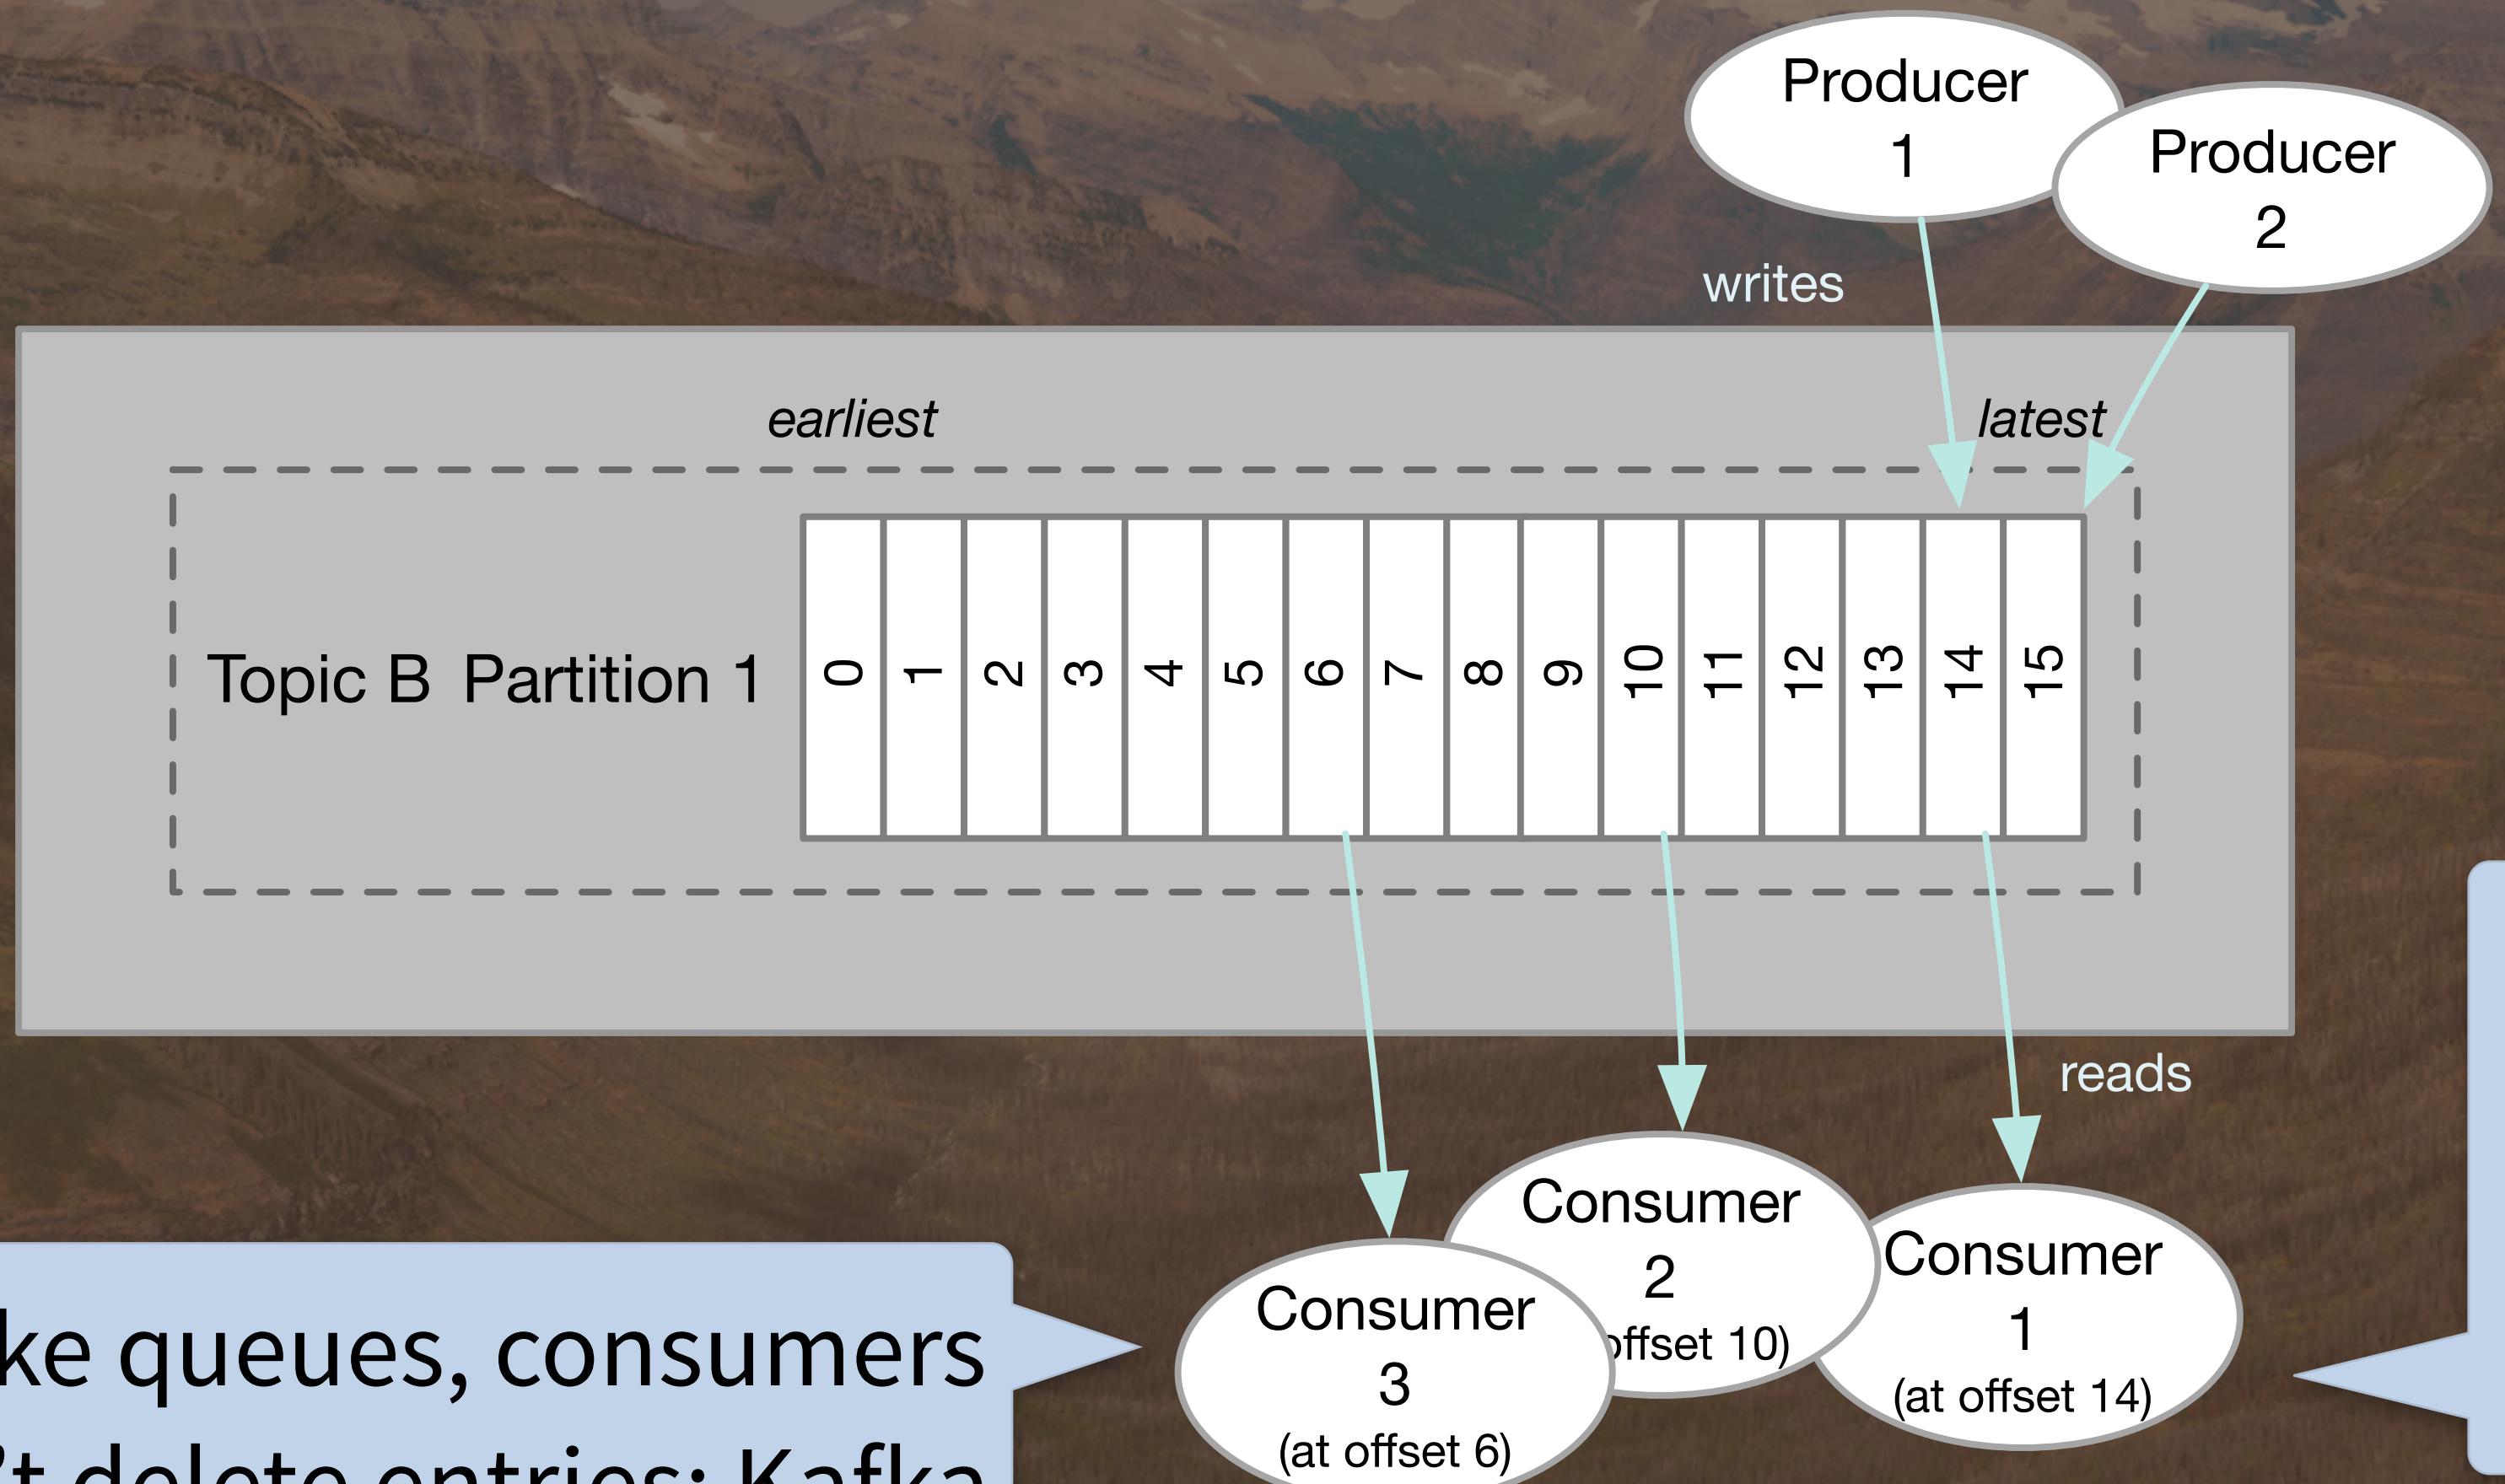


Organized into topics

Topics are partitioned,  
replicated, and  
distributed



Logs, not queues!

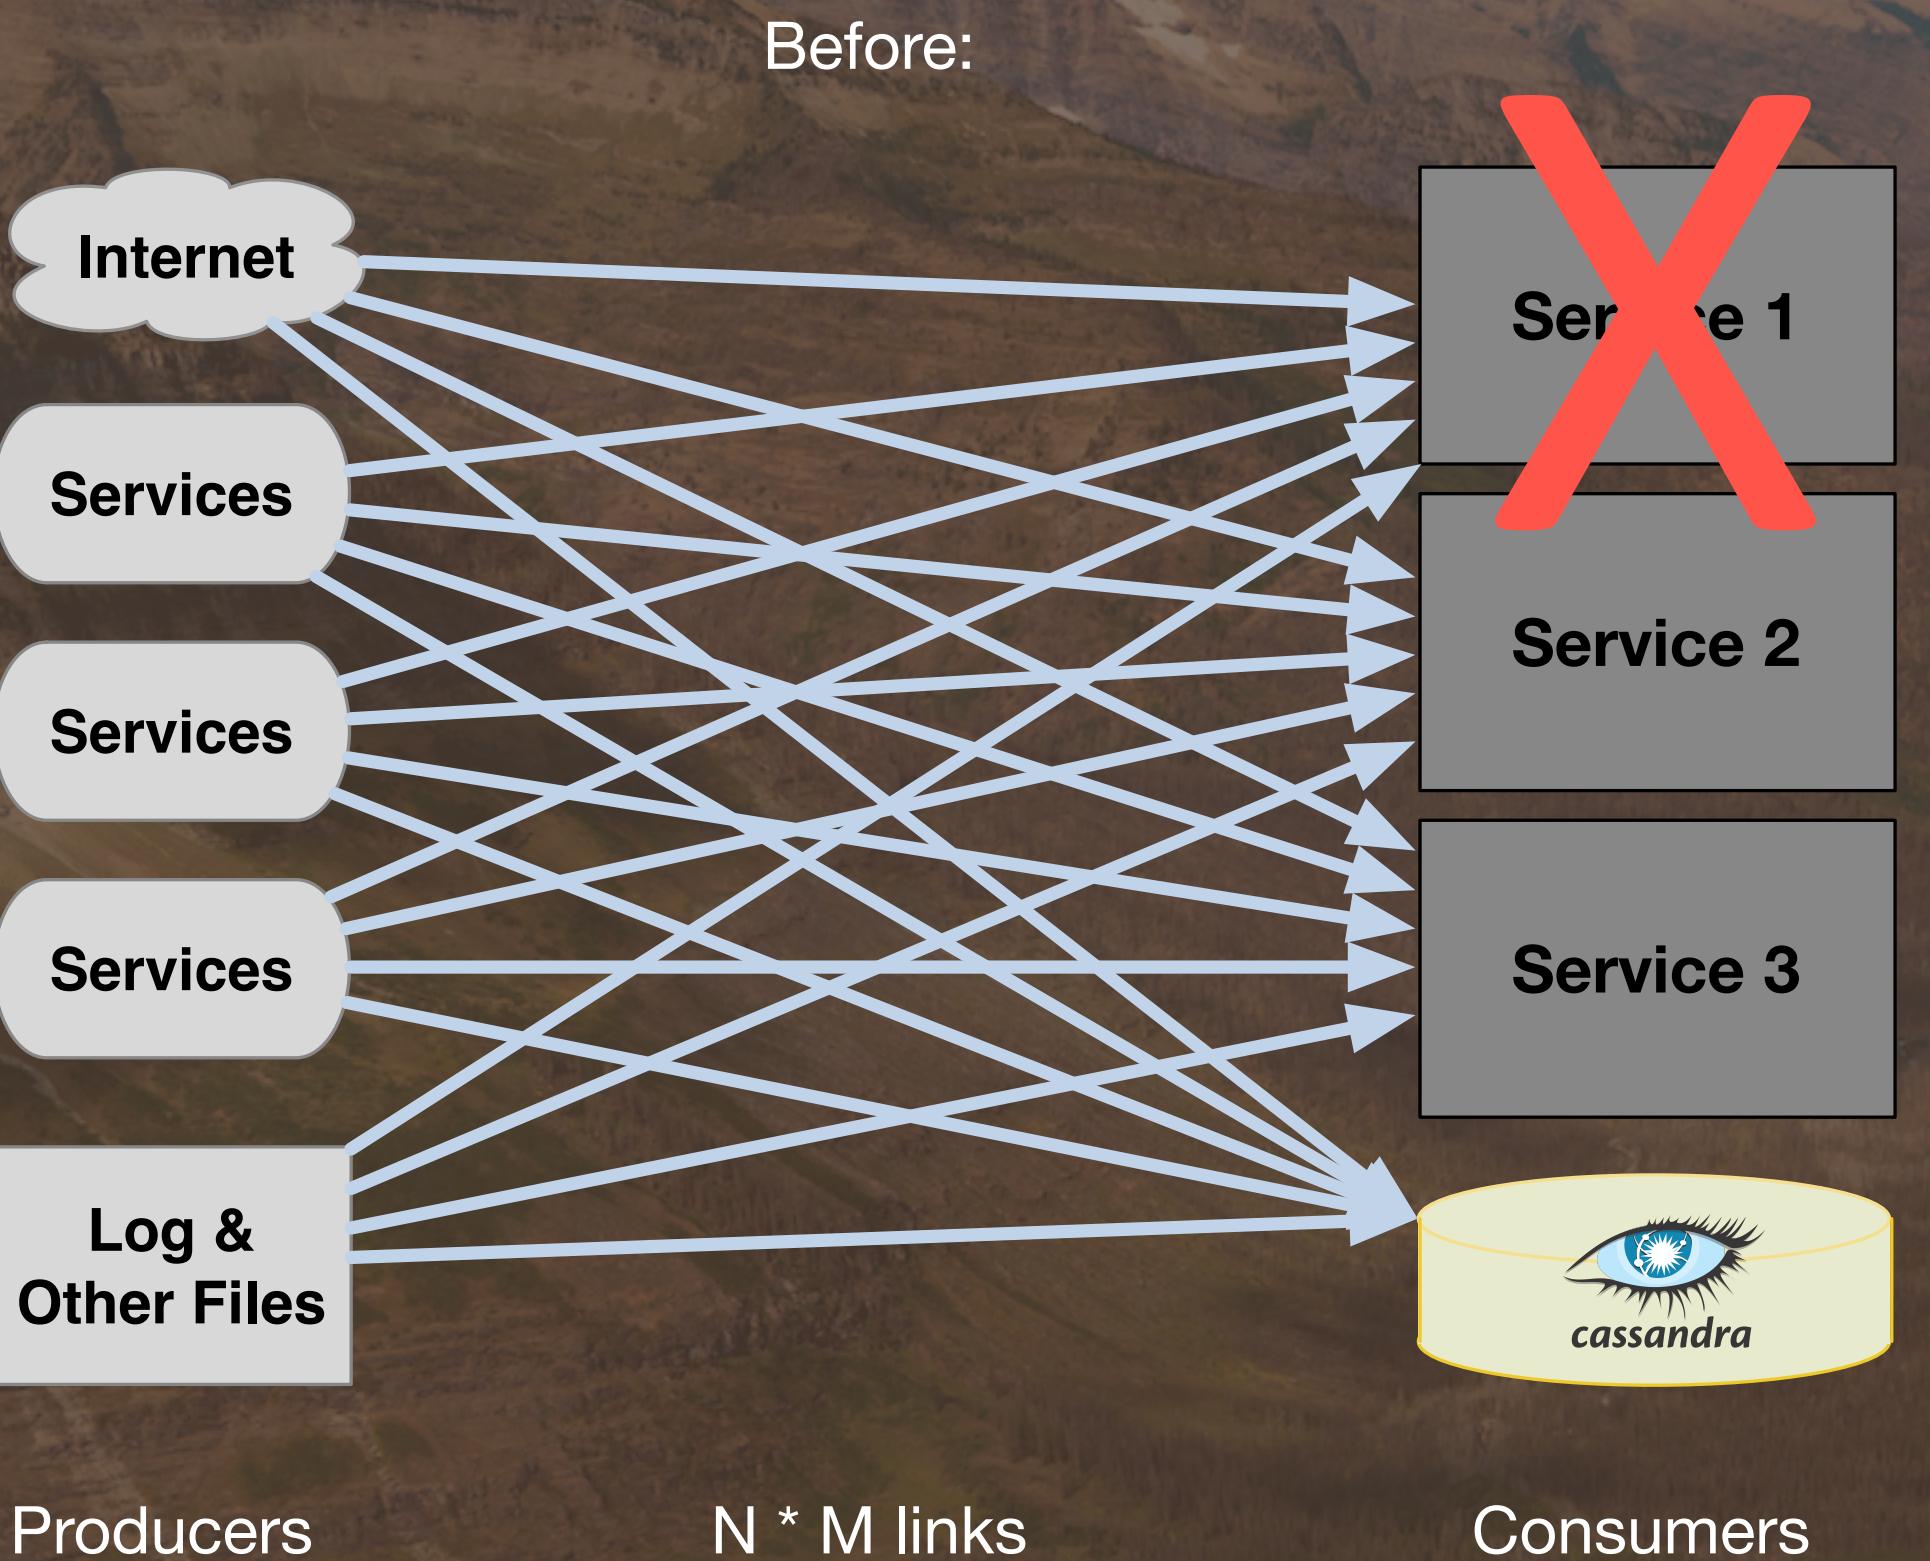


Unlike queues, consumers don't delete entries; Kafka manages their lifecycles

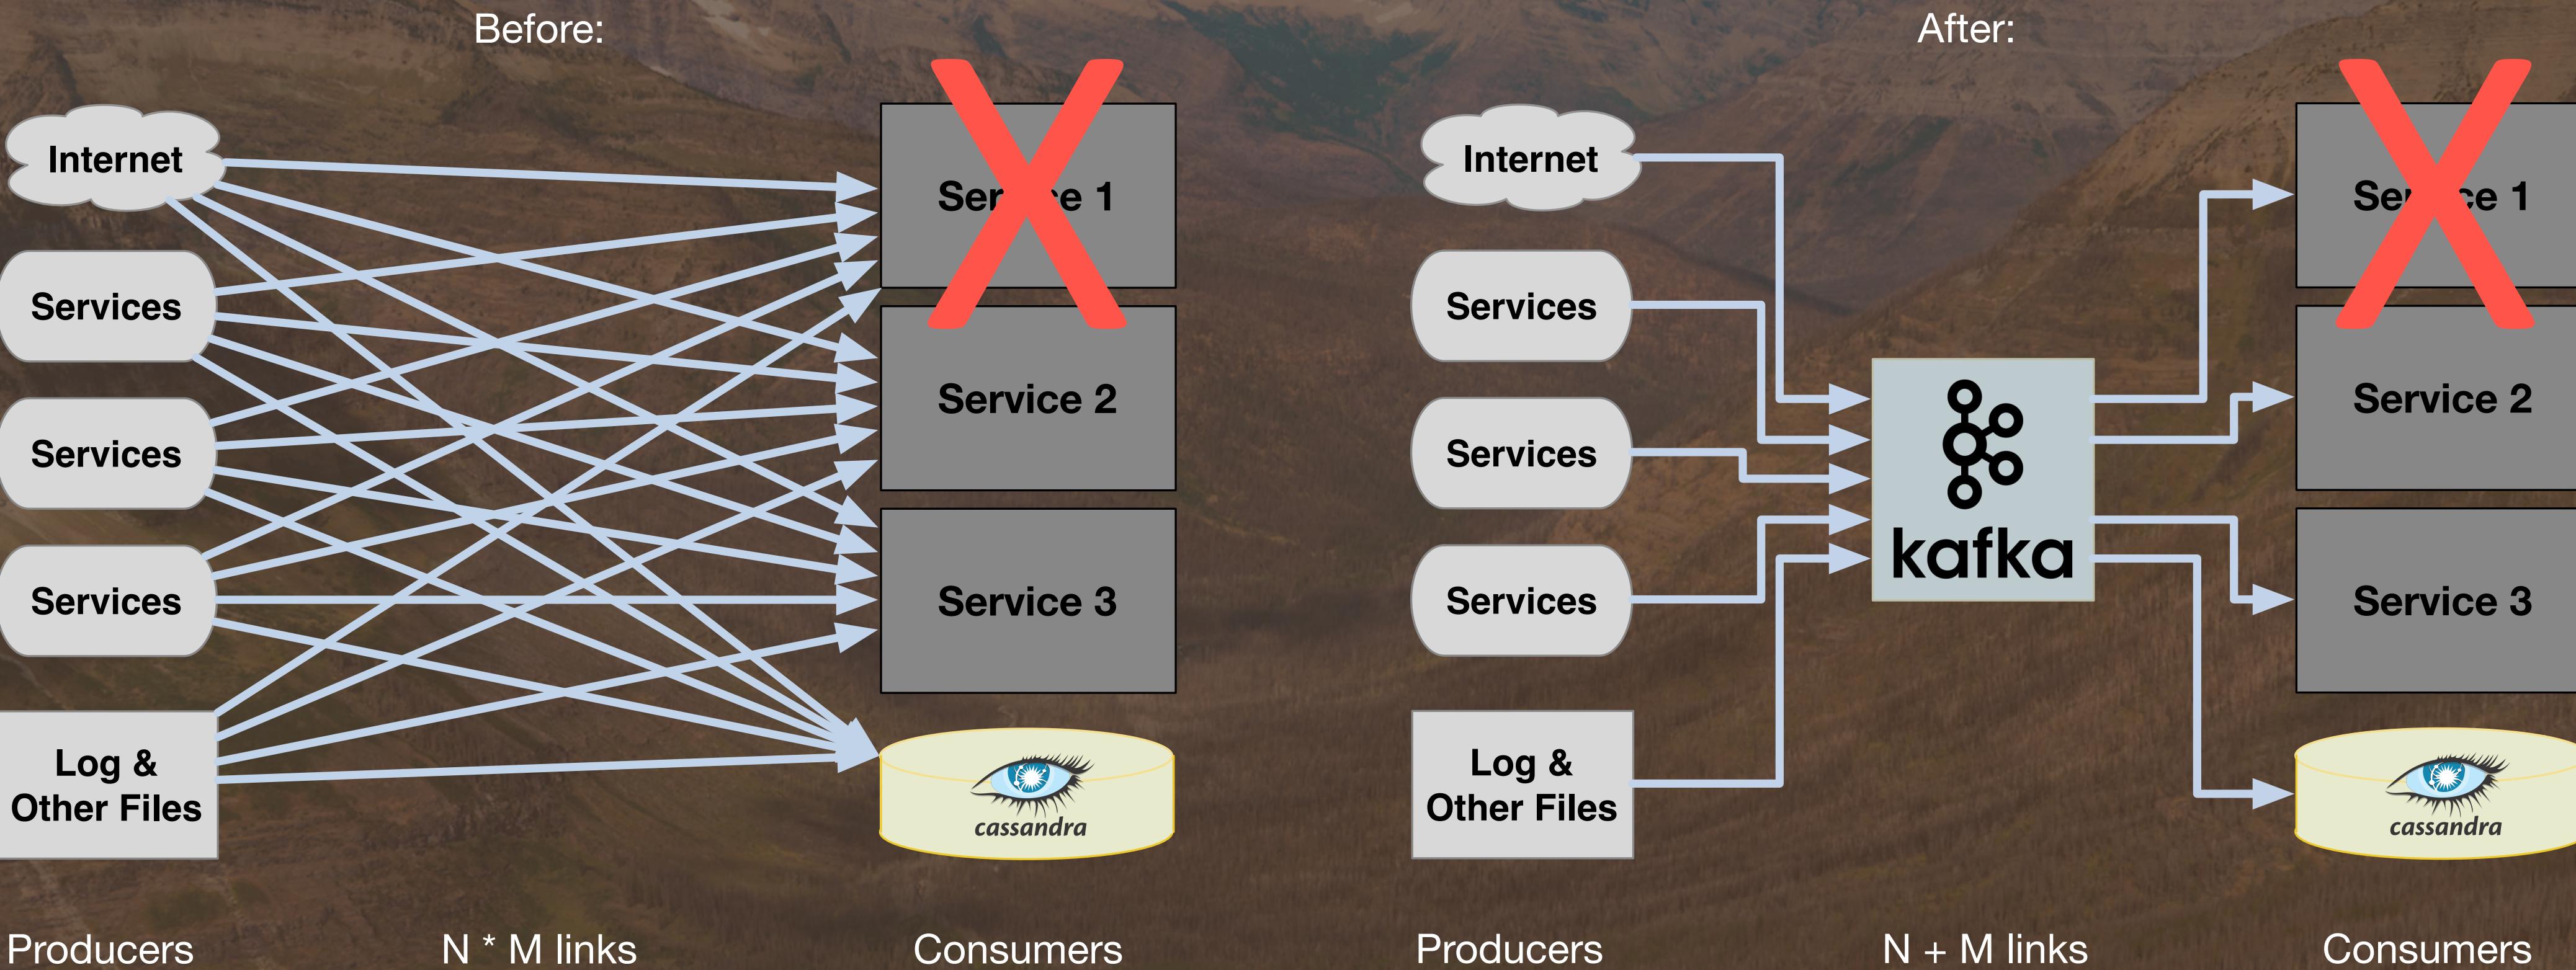
M Producers

N Consumers,  
who start  
reading where  
they want

# Architectural Benefits of Kafka

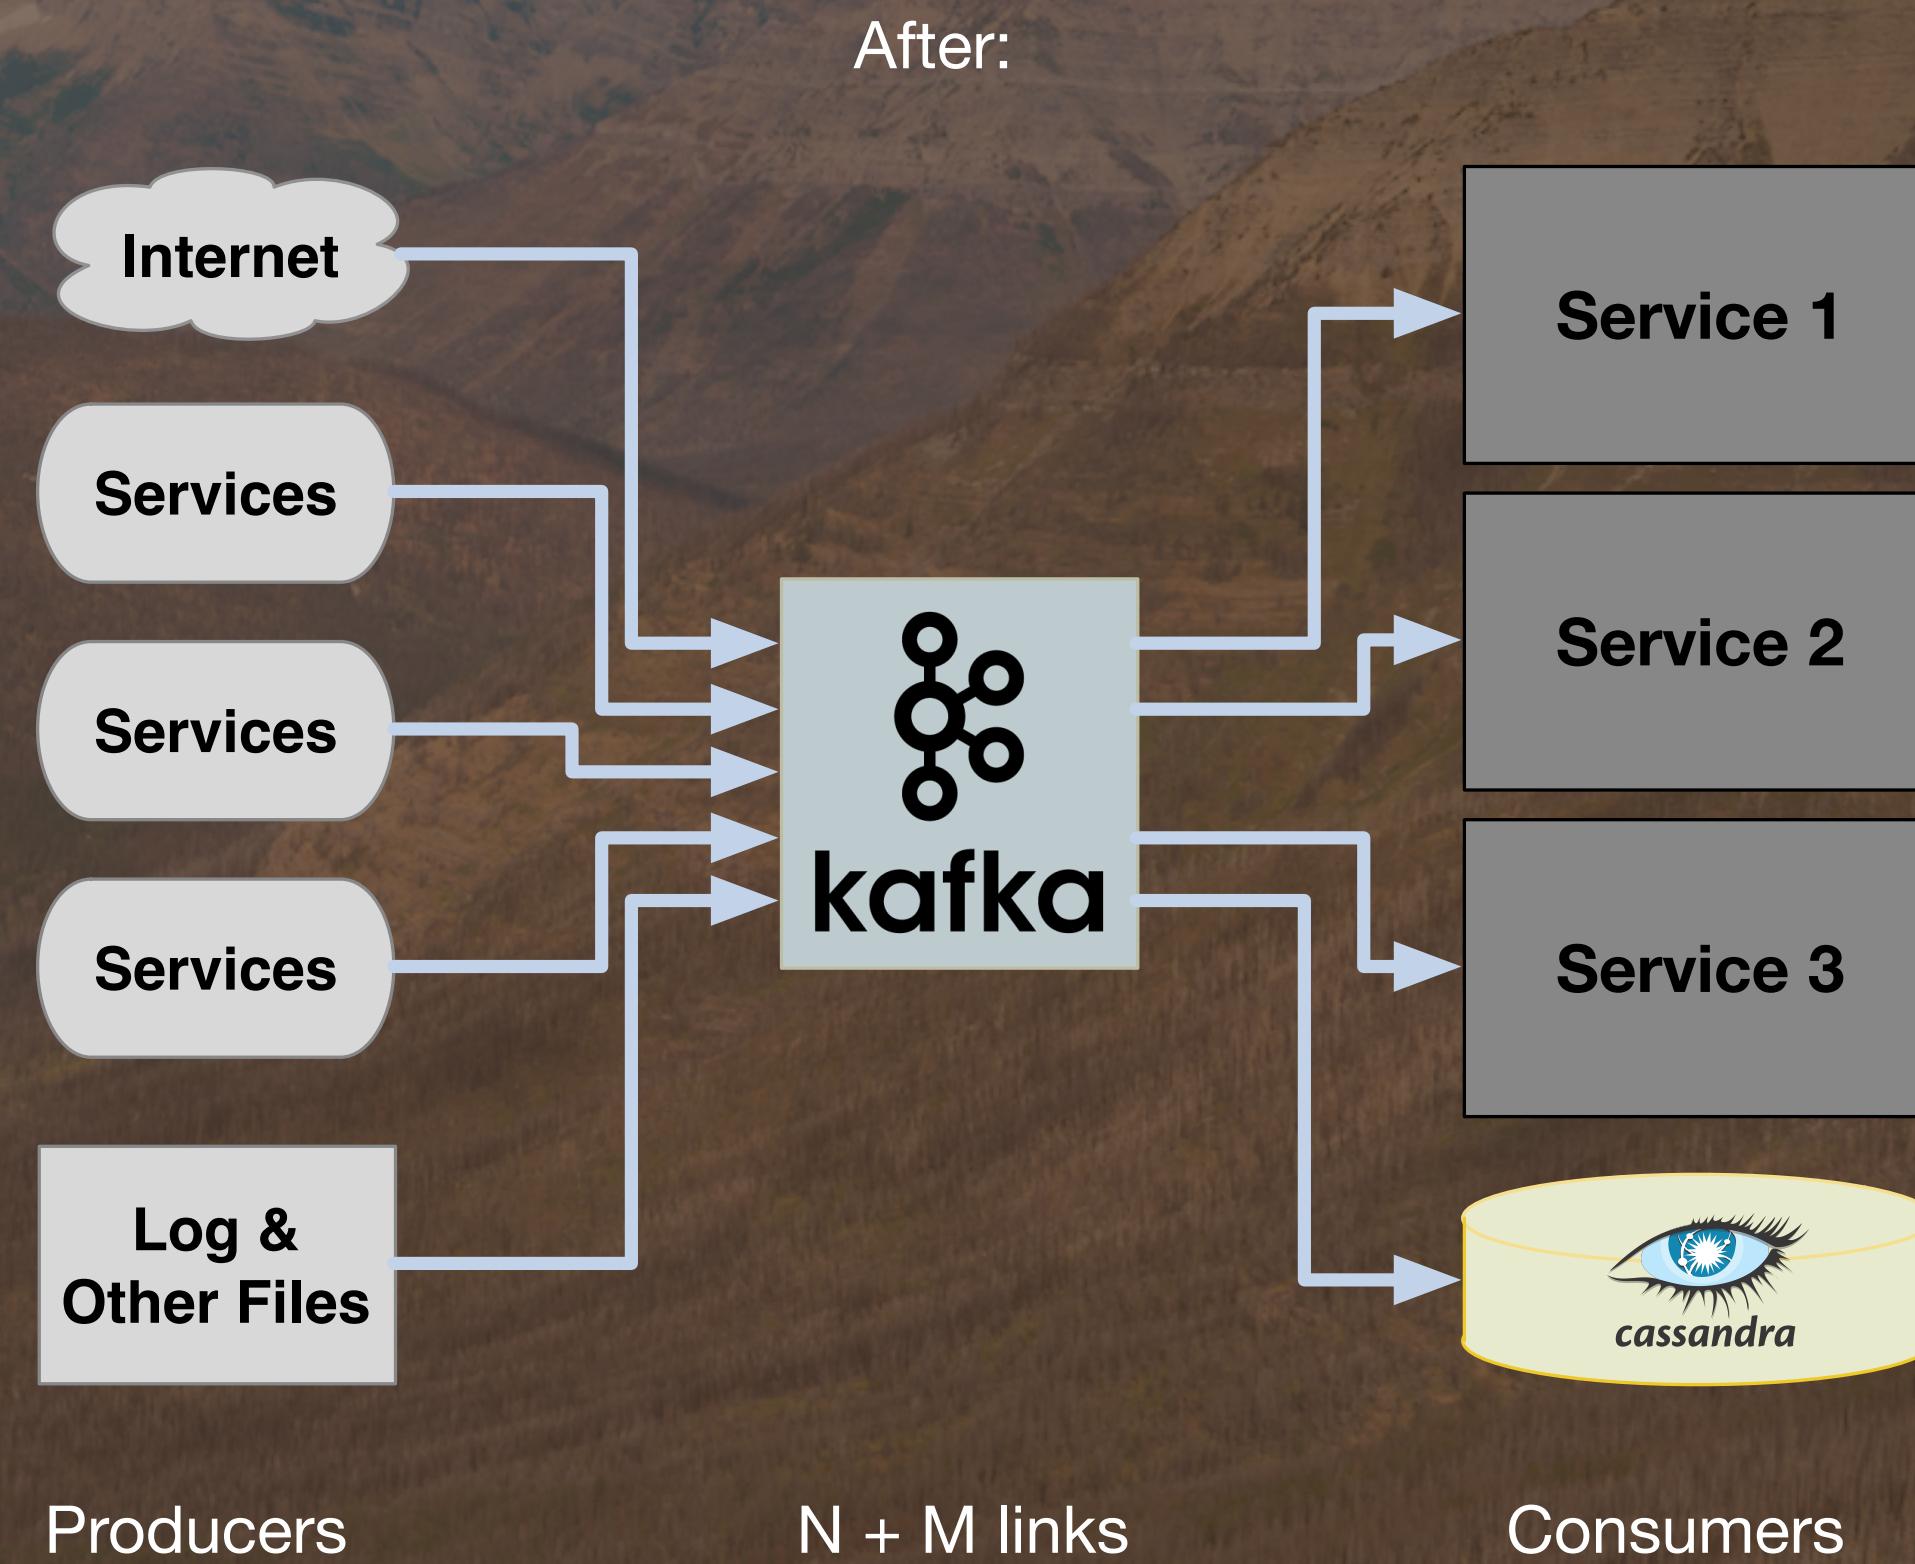


# Architectural Benefits of Kafka



# Architectural Benefits of Kafka

- Simplify dependencies
- Resilient against data loss
- M producers, N consumers
- Simplicity of one “API” for communication



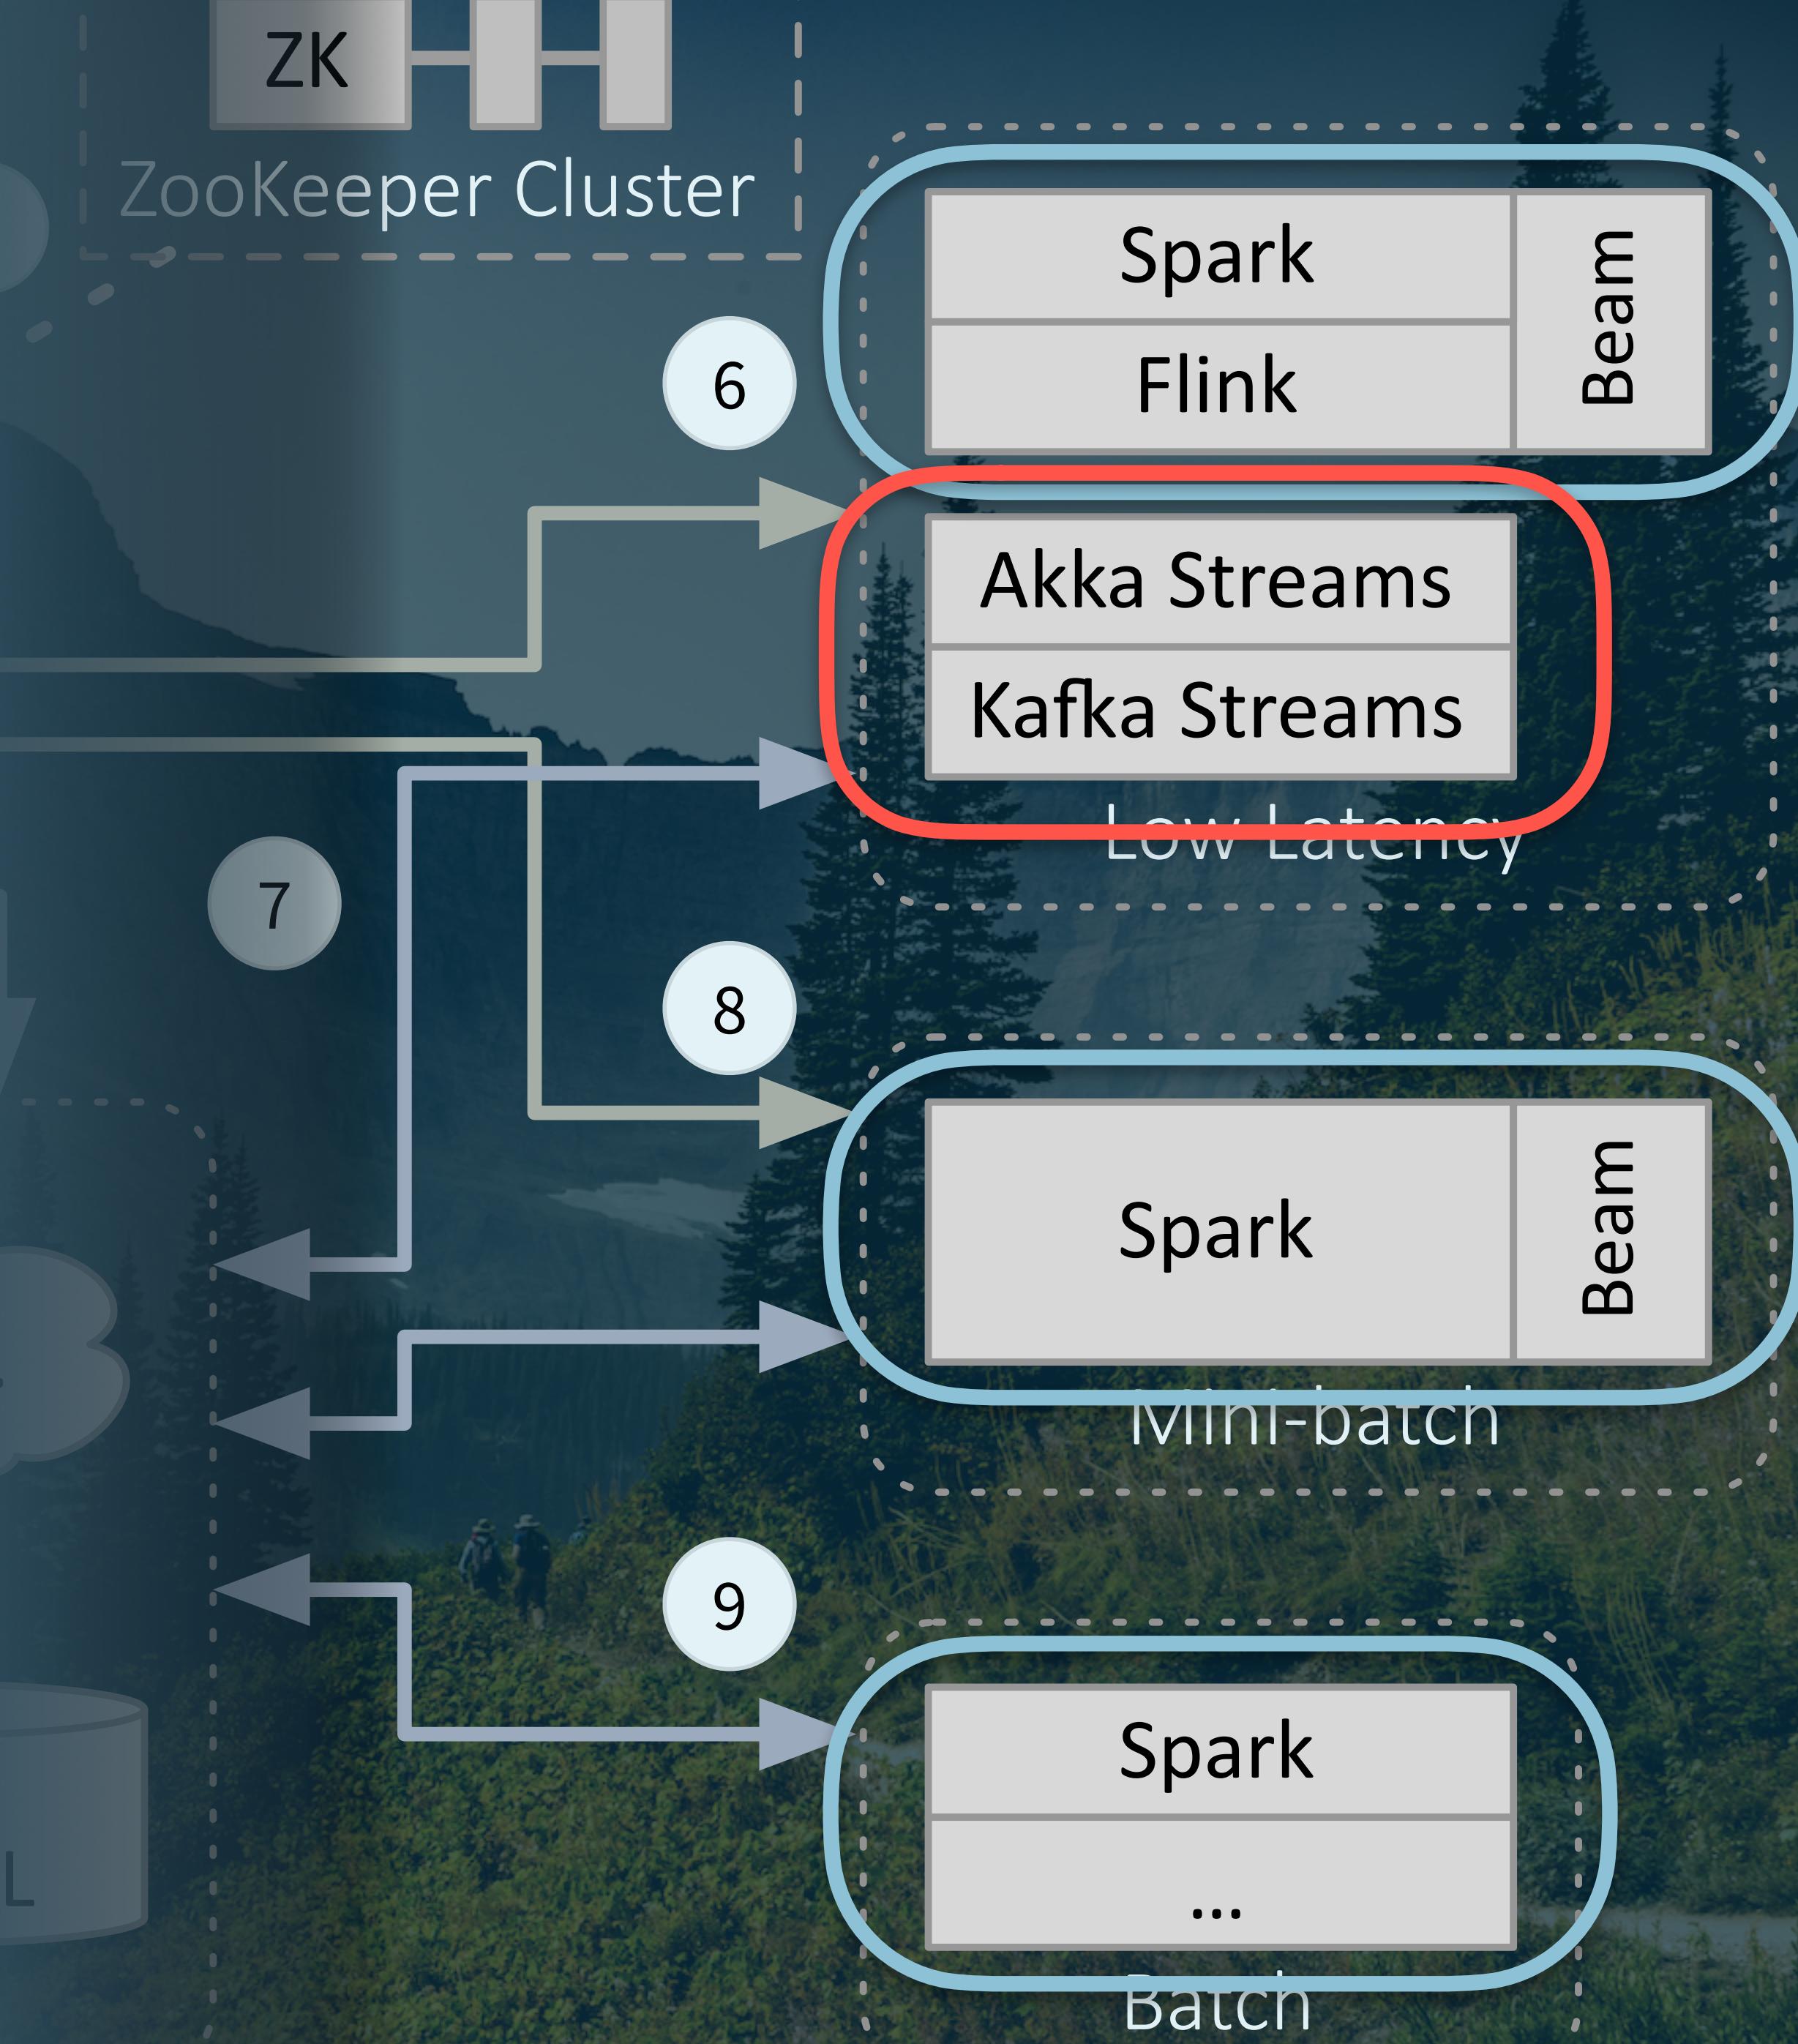
# Streaming Engines



# Streaming Engines

Should you use a scalable data system like Spark or Flink?

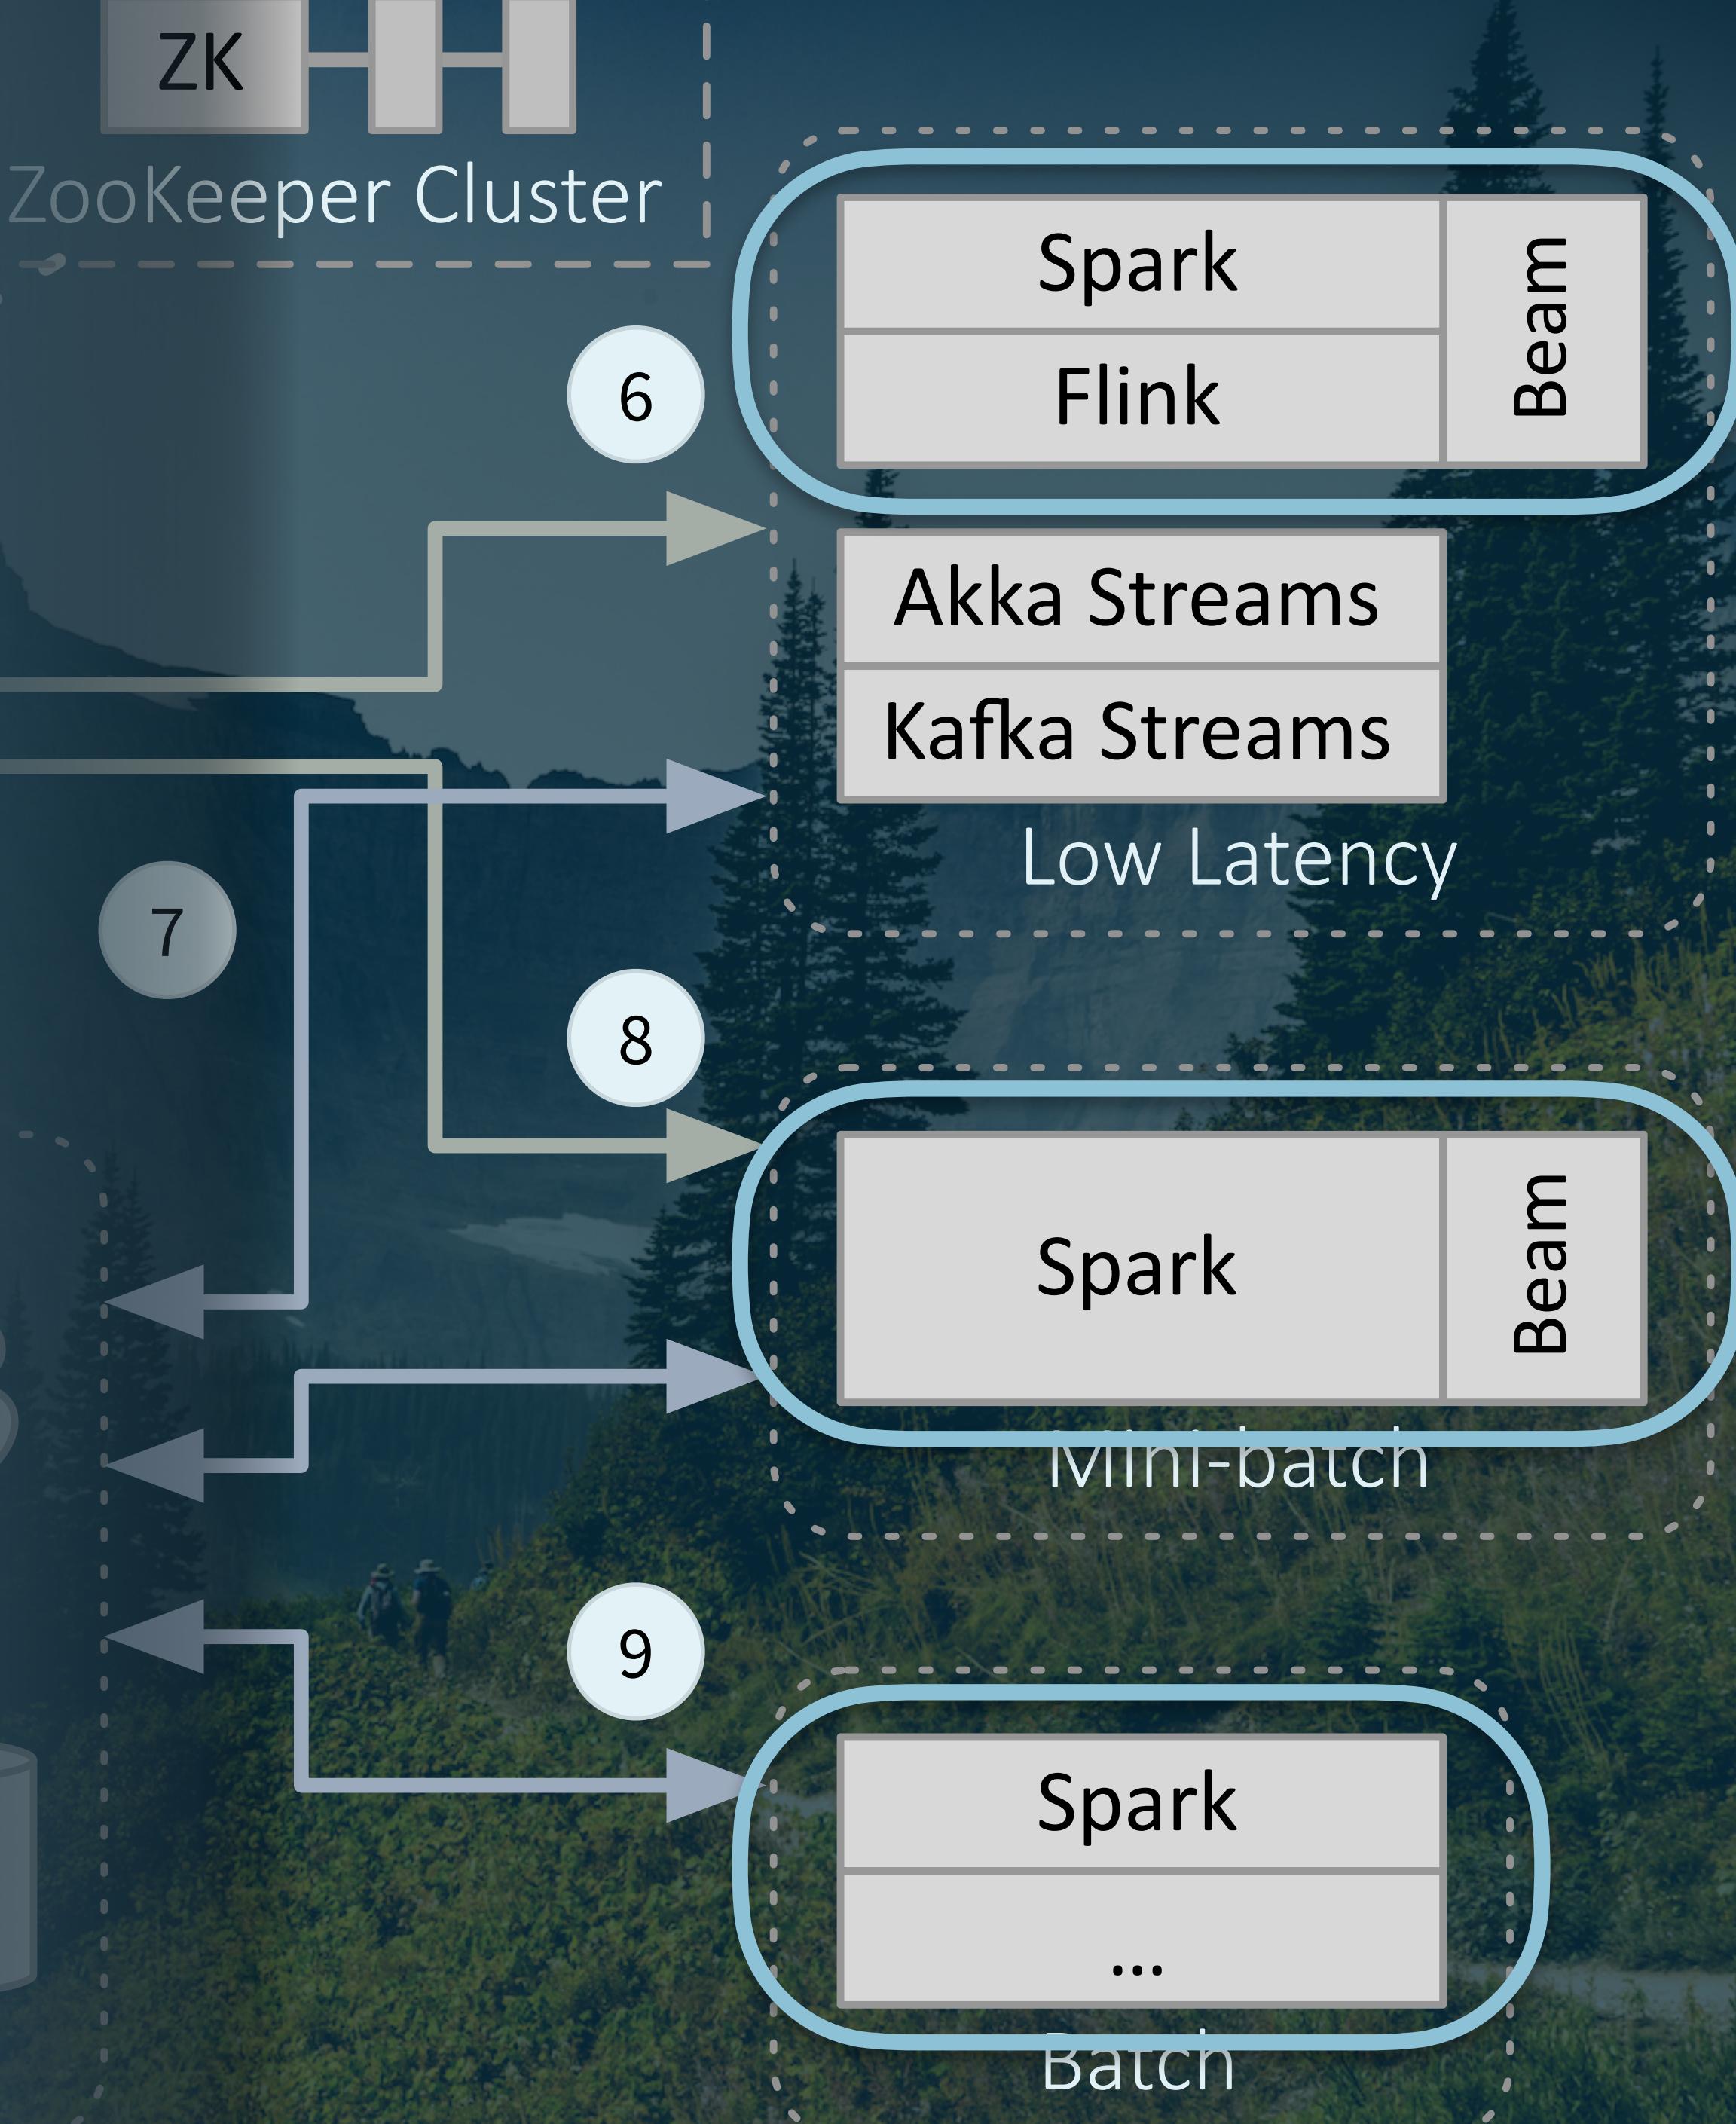
Or should you use a *library* in “classic” microservices?



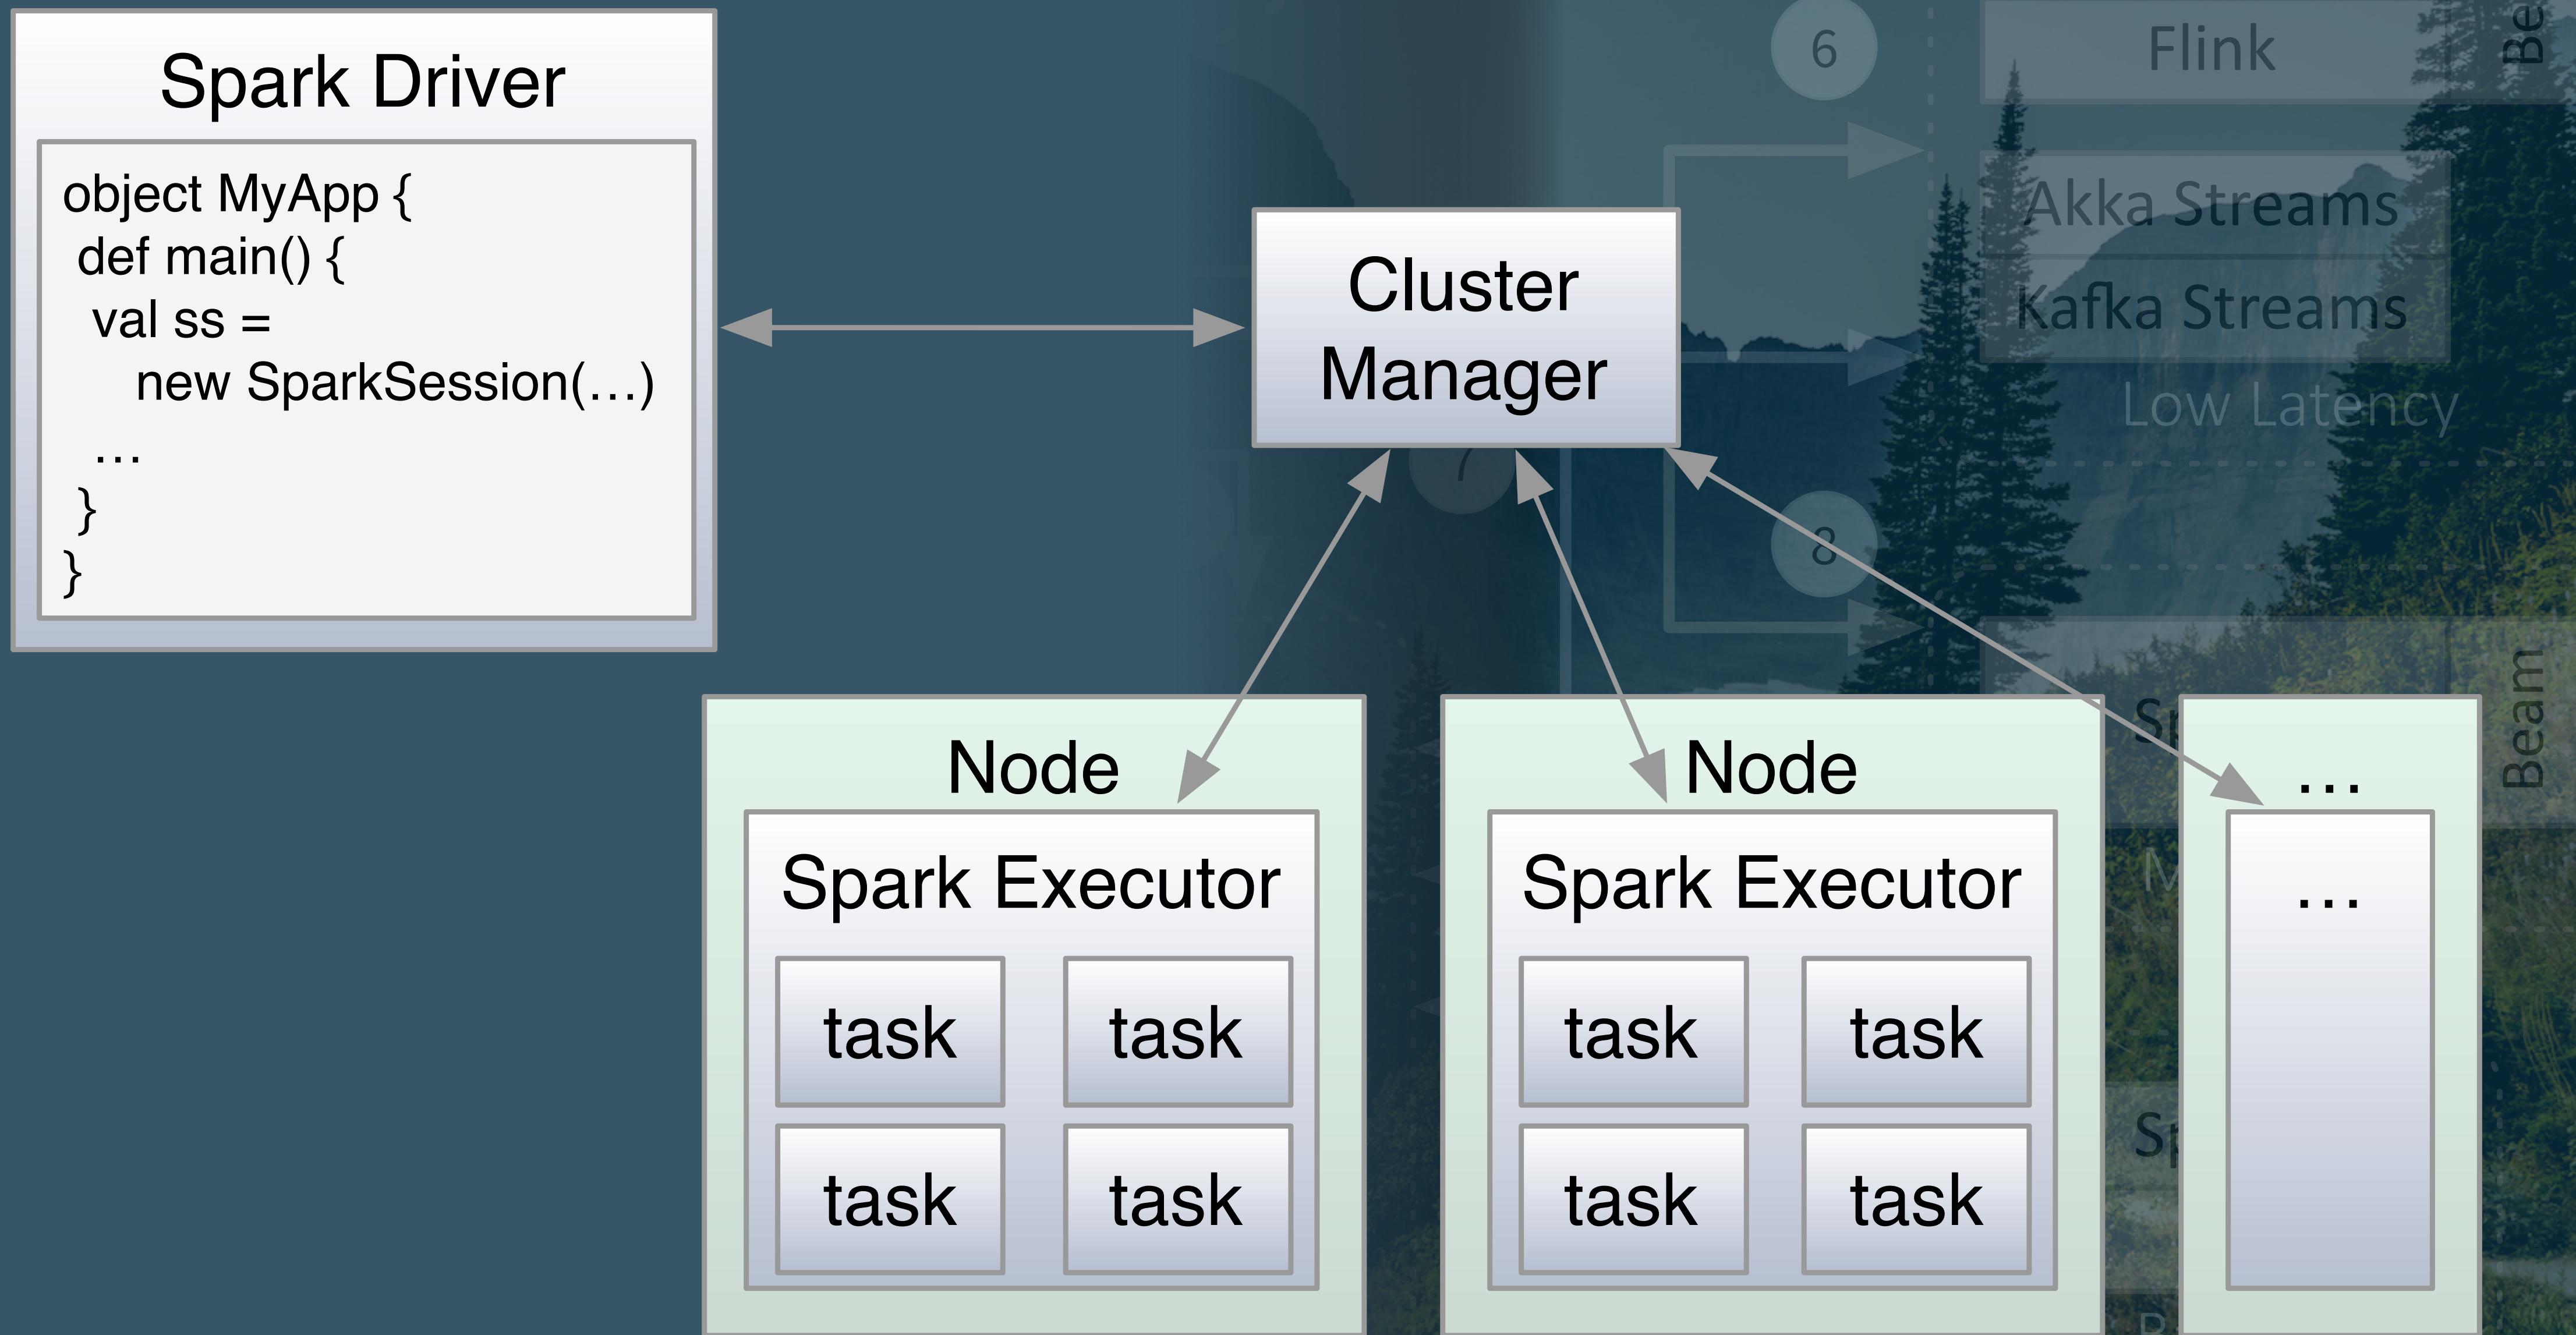
# Streaming Engines

Spark, Flink - services to which you submit work. Large scale, automatic data partitioning.

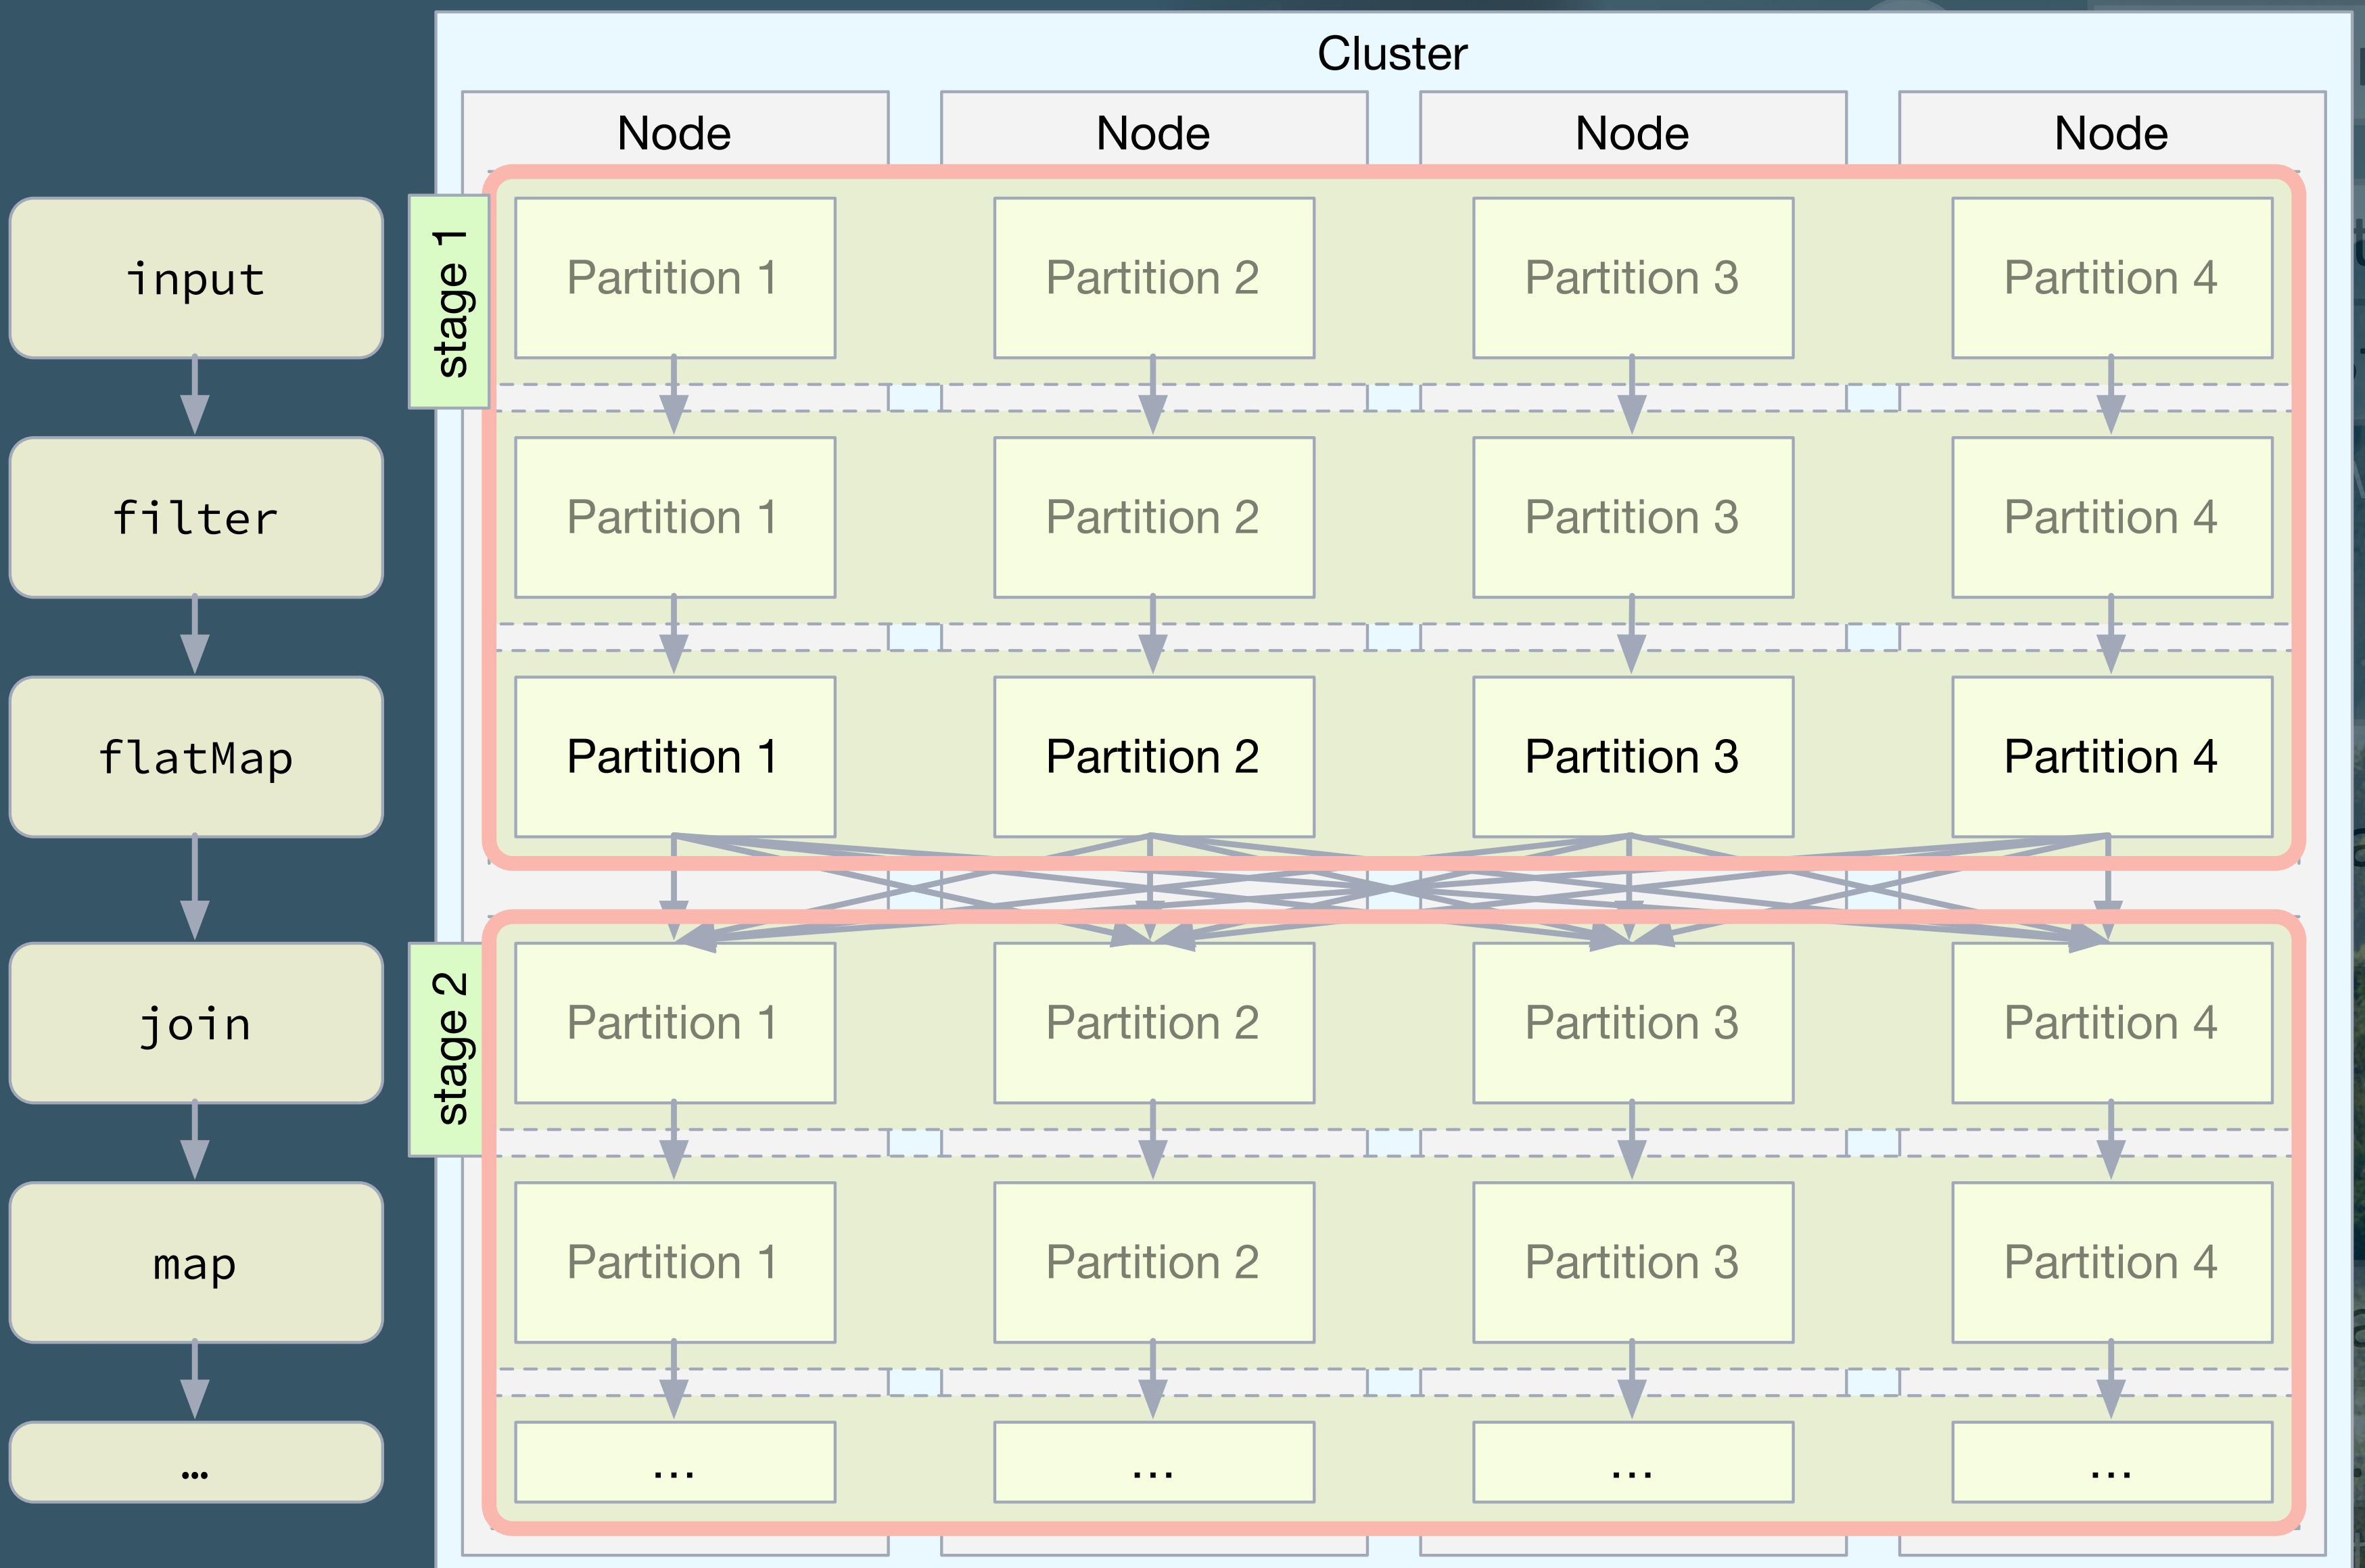
Beam - similar. Google's project that has been instrumental in defining streaming semantics.



# They do a *lot* (Spark example)



# They do a lot (Spark example)



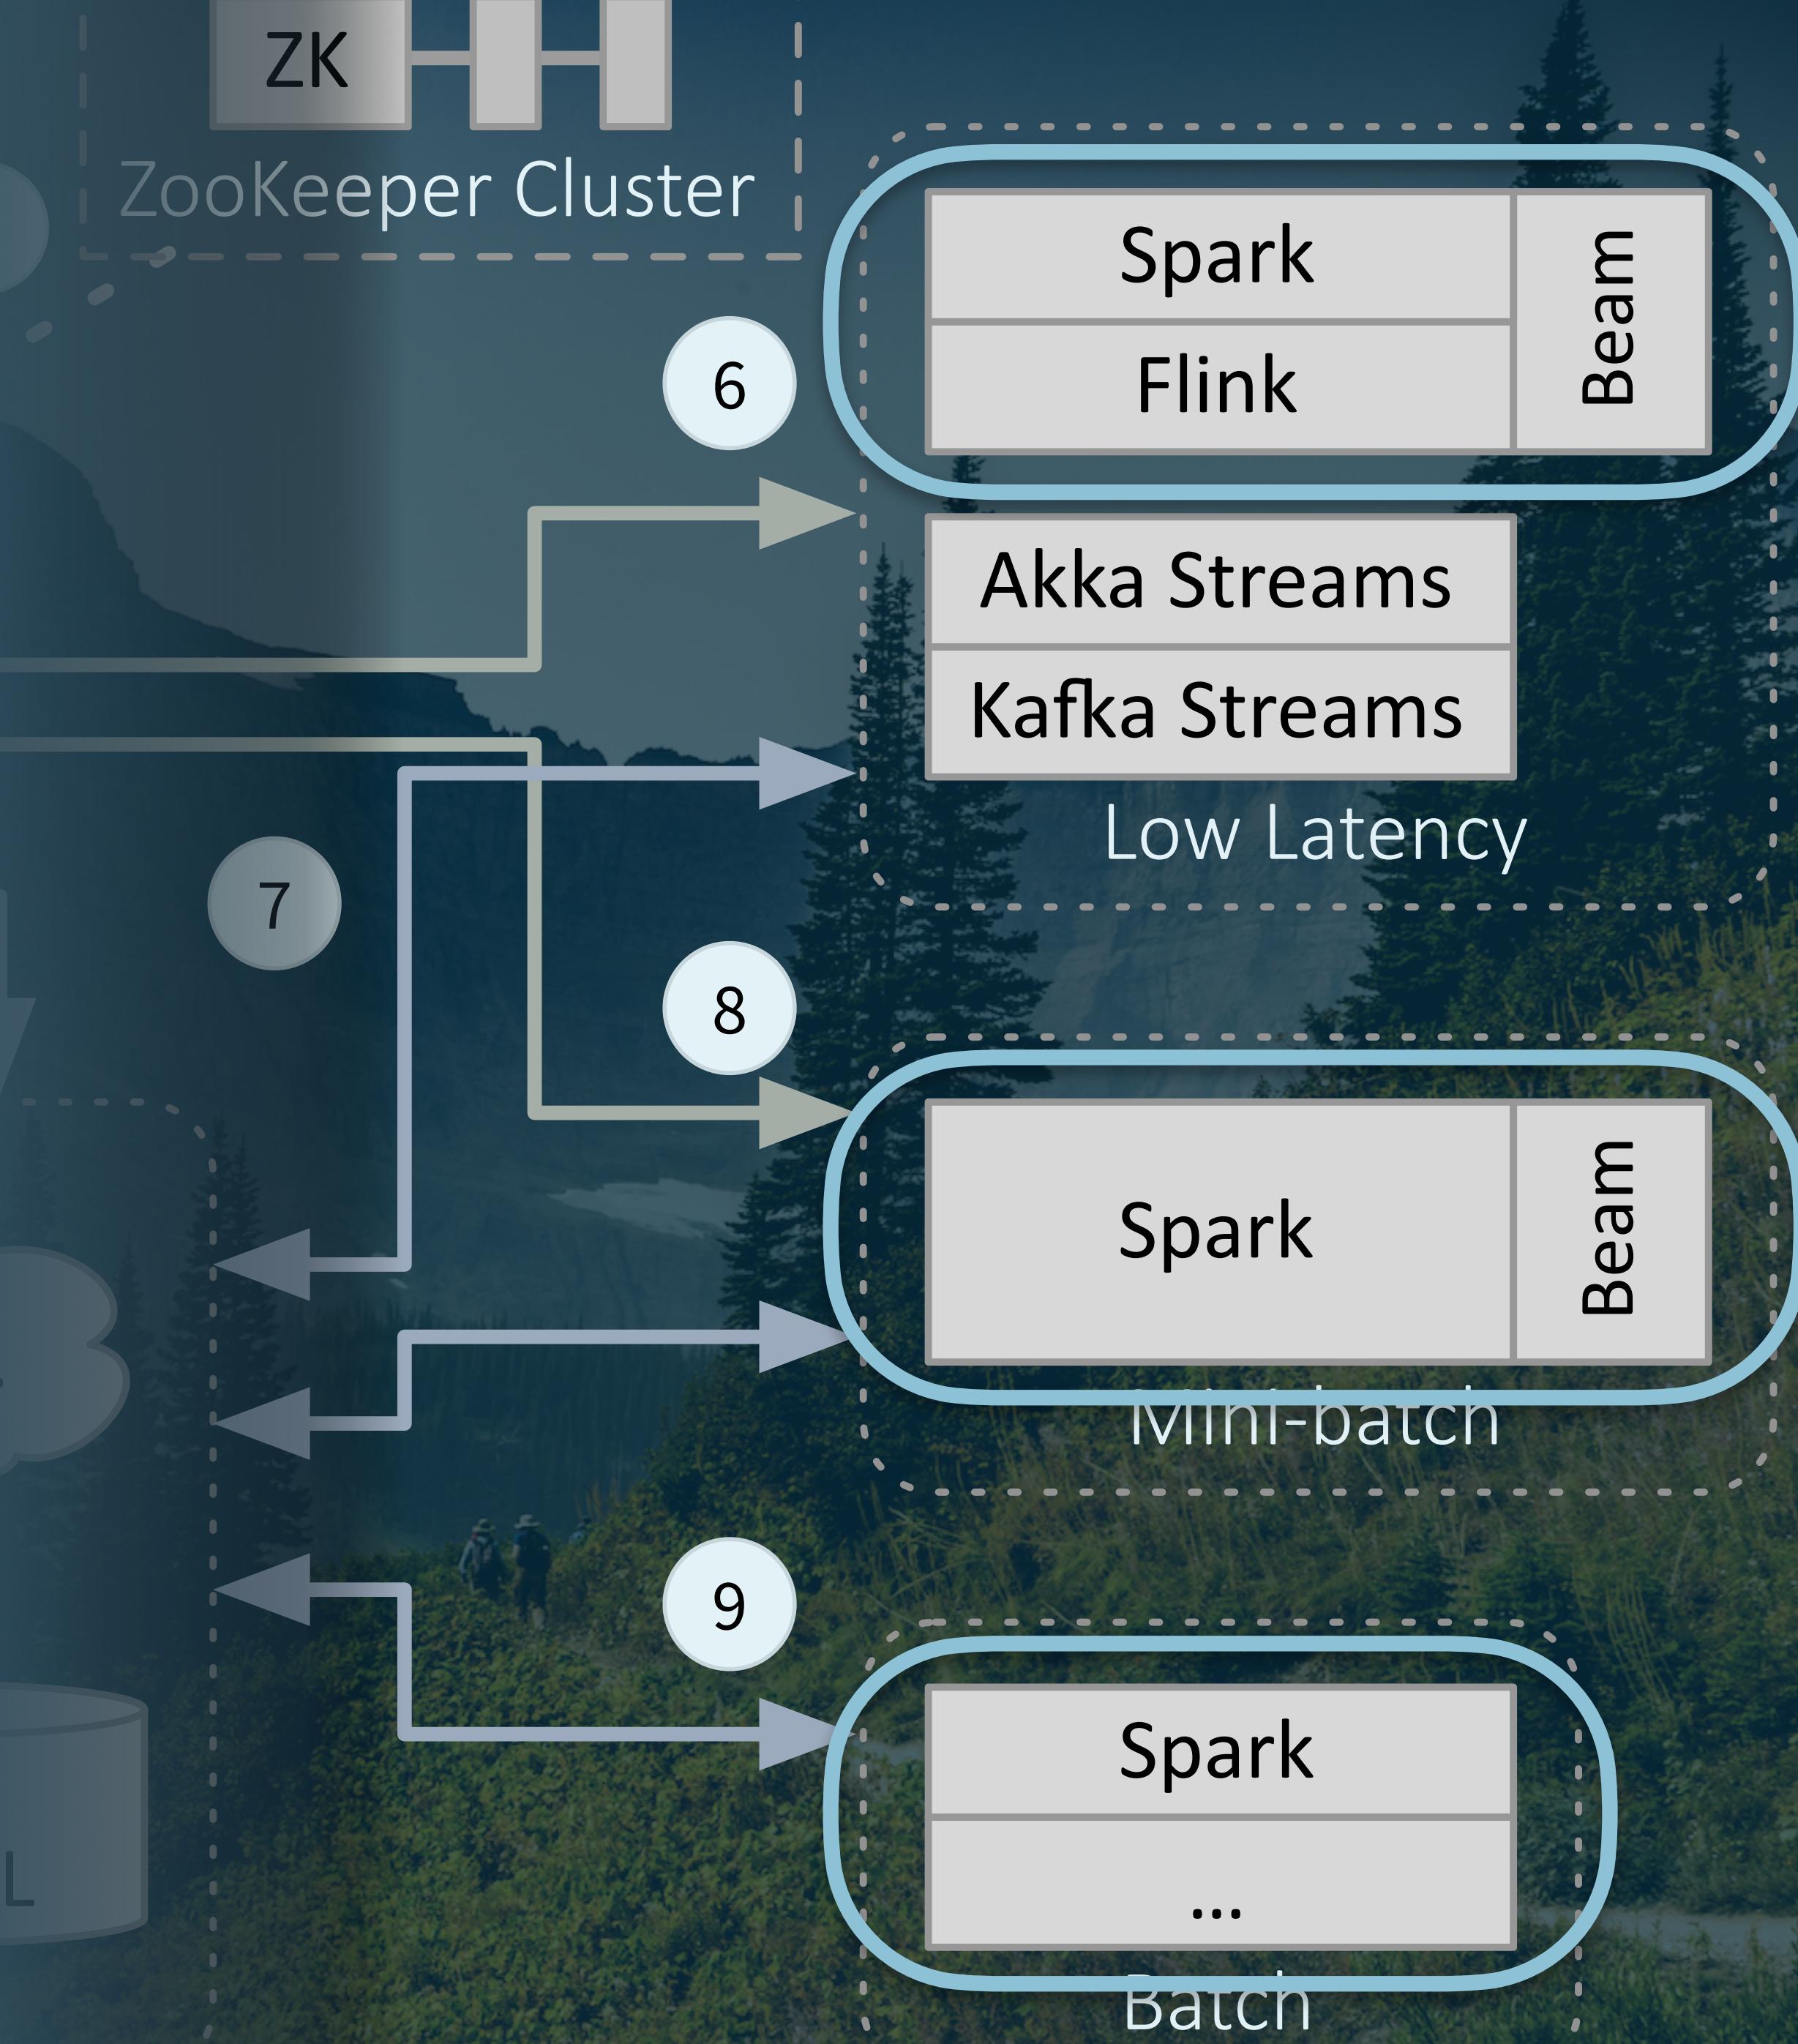
# Spark and Flink

Wired:

- Massive scalability
- Rich semantics

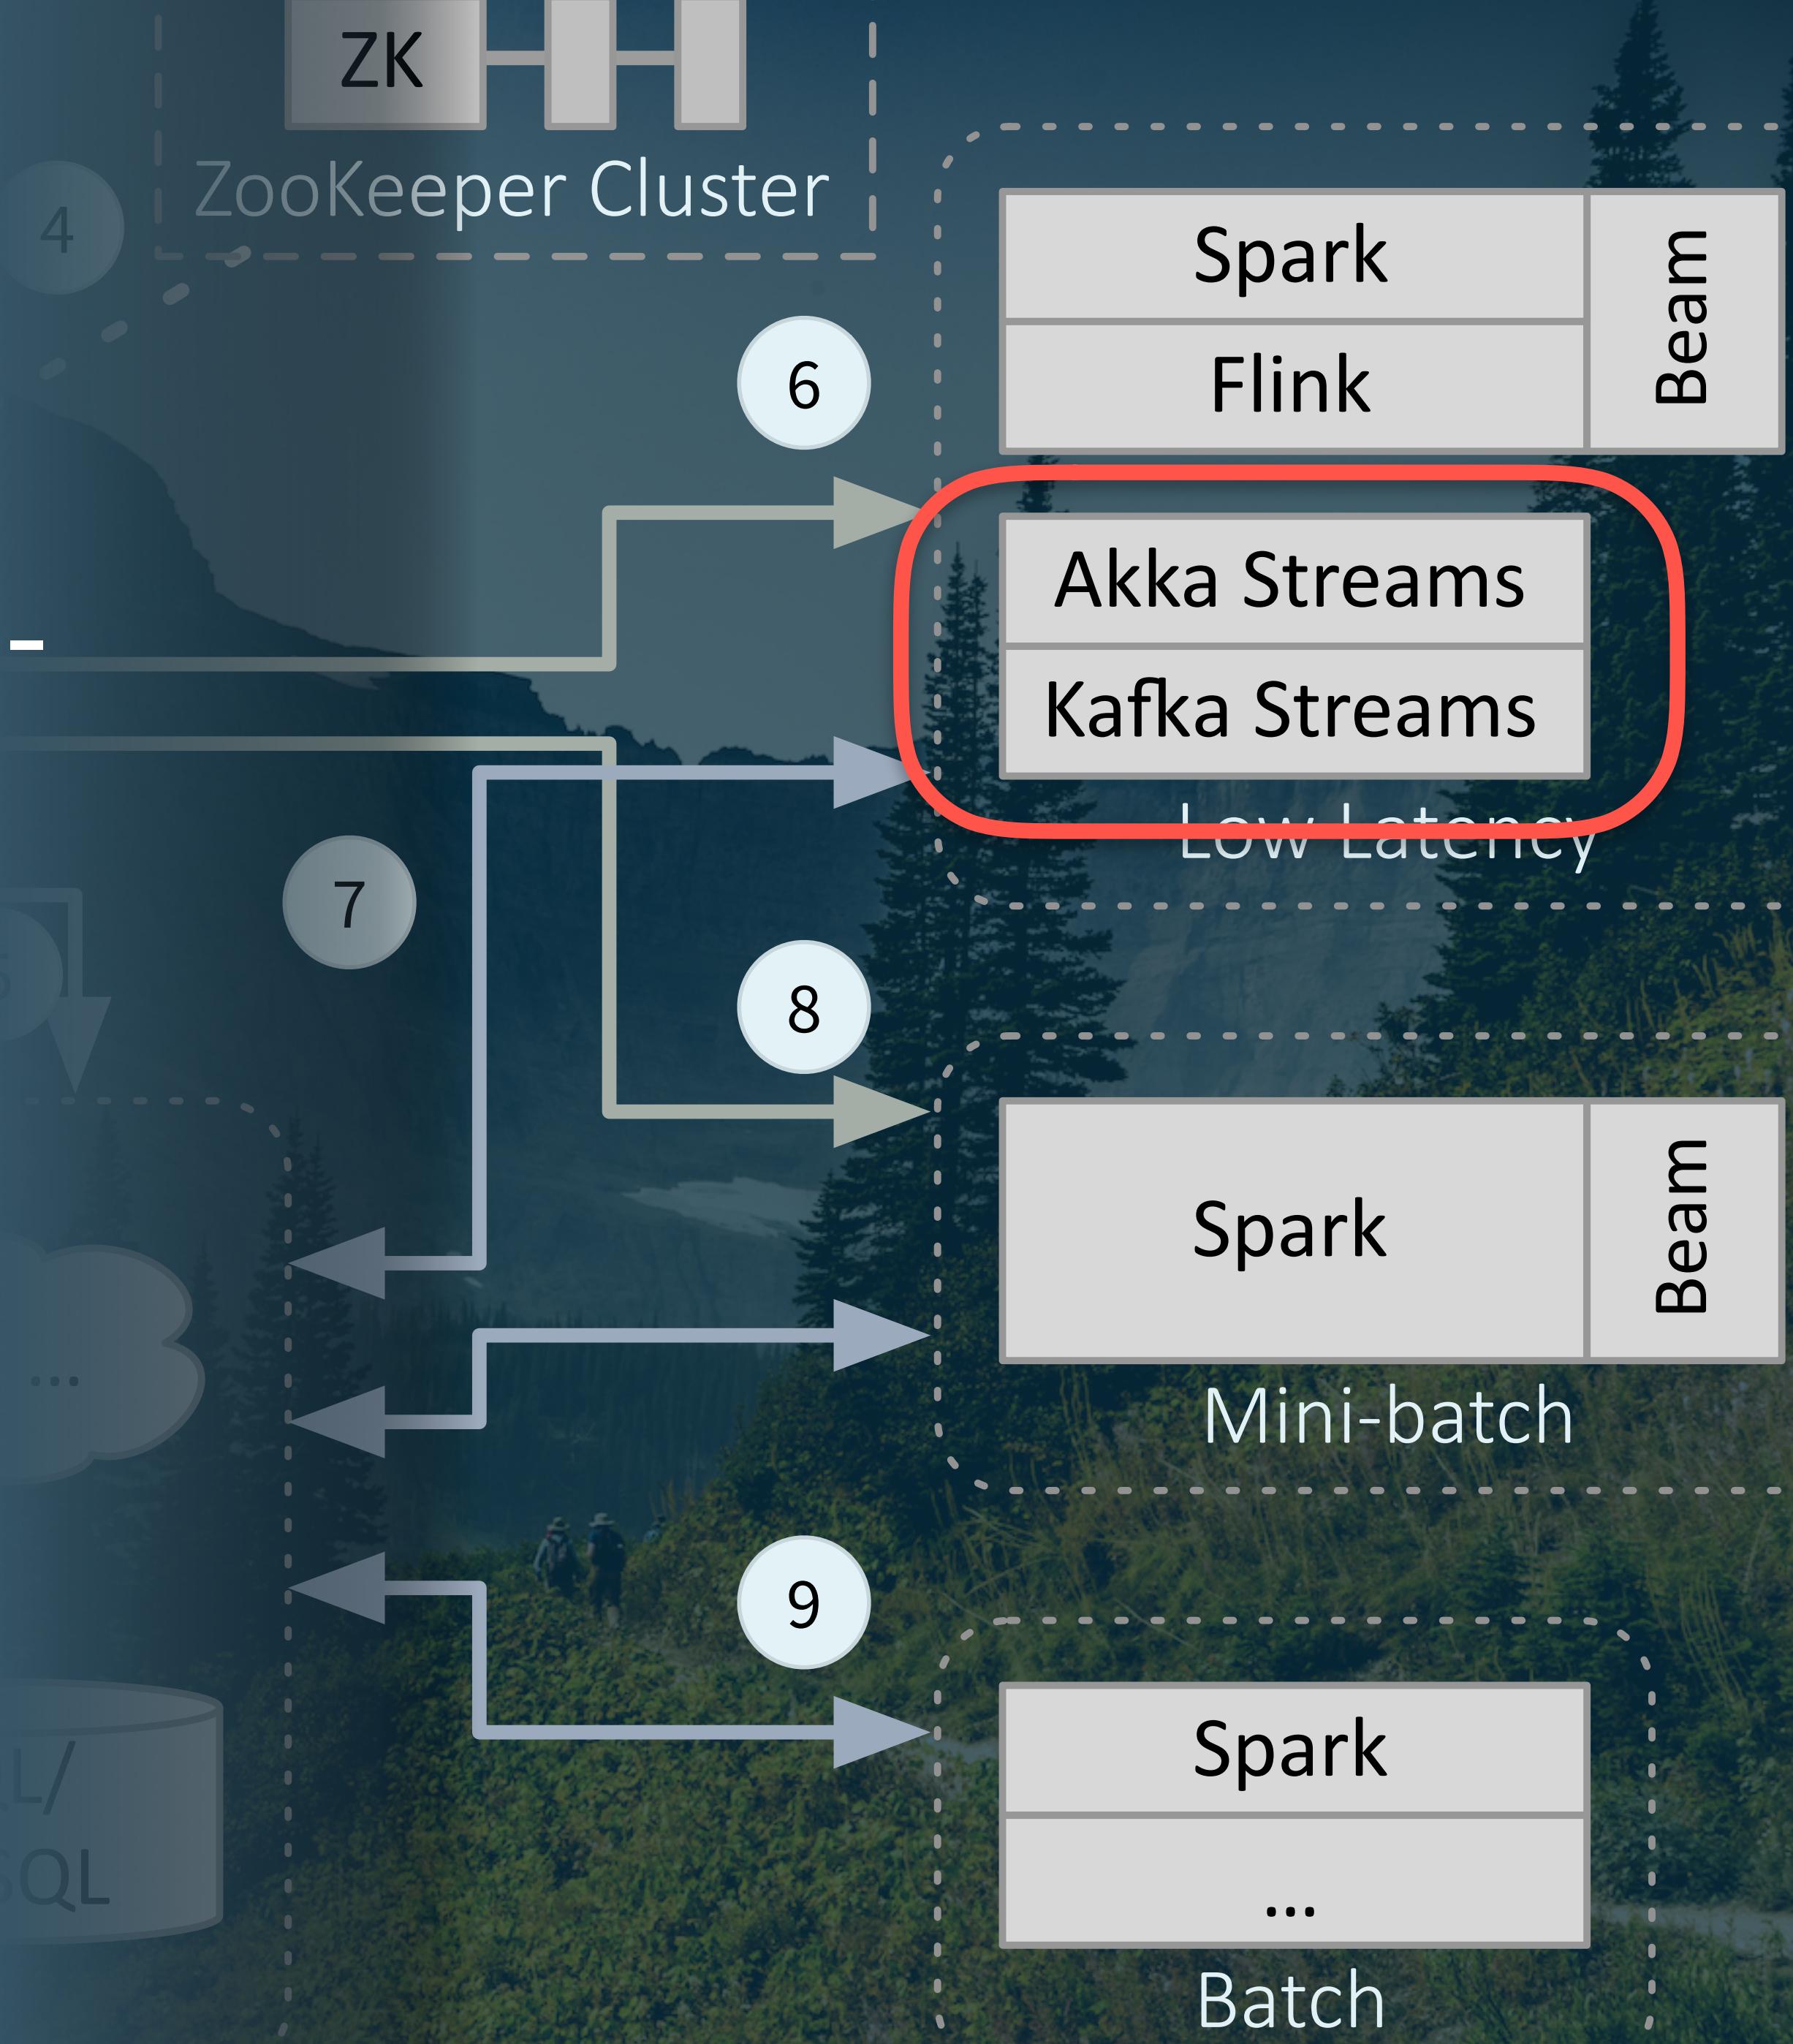
Tired:

- Latency
- Integration in CI/CD workflows



# Streaming Engines

Akka Streams, Kafka Streams - libraries for “data-centric microservices”. Smaller scale, but great flexibility



# Microservice All the Things!!



**Scott Hanselman**

@shanselman

Follow

Microservices, for when your in-process methods have too little latency.

**Dave Cheney** @davecheney

Microservices, for when function calls are too reliable.

4:11 AM - 25 Feb 2018

**207** Retweets **566** Likes



25

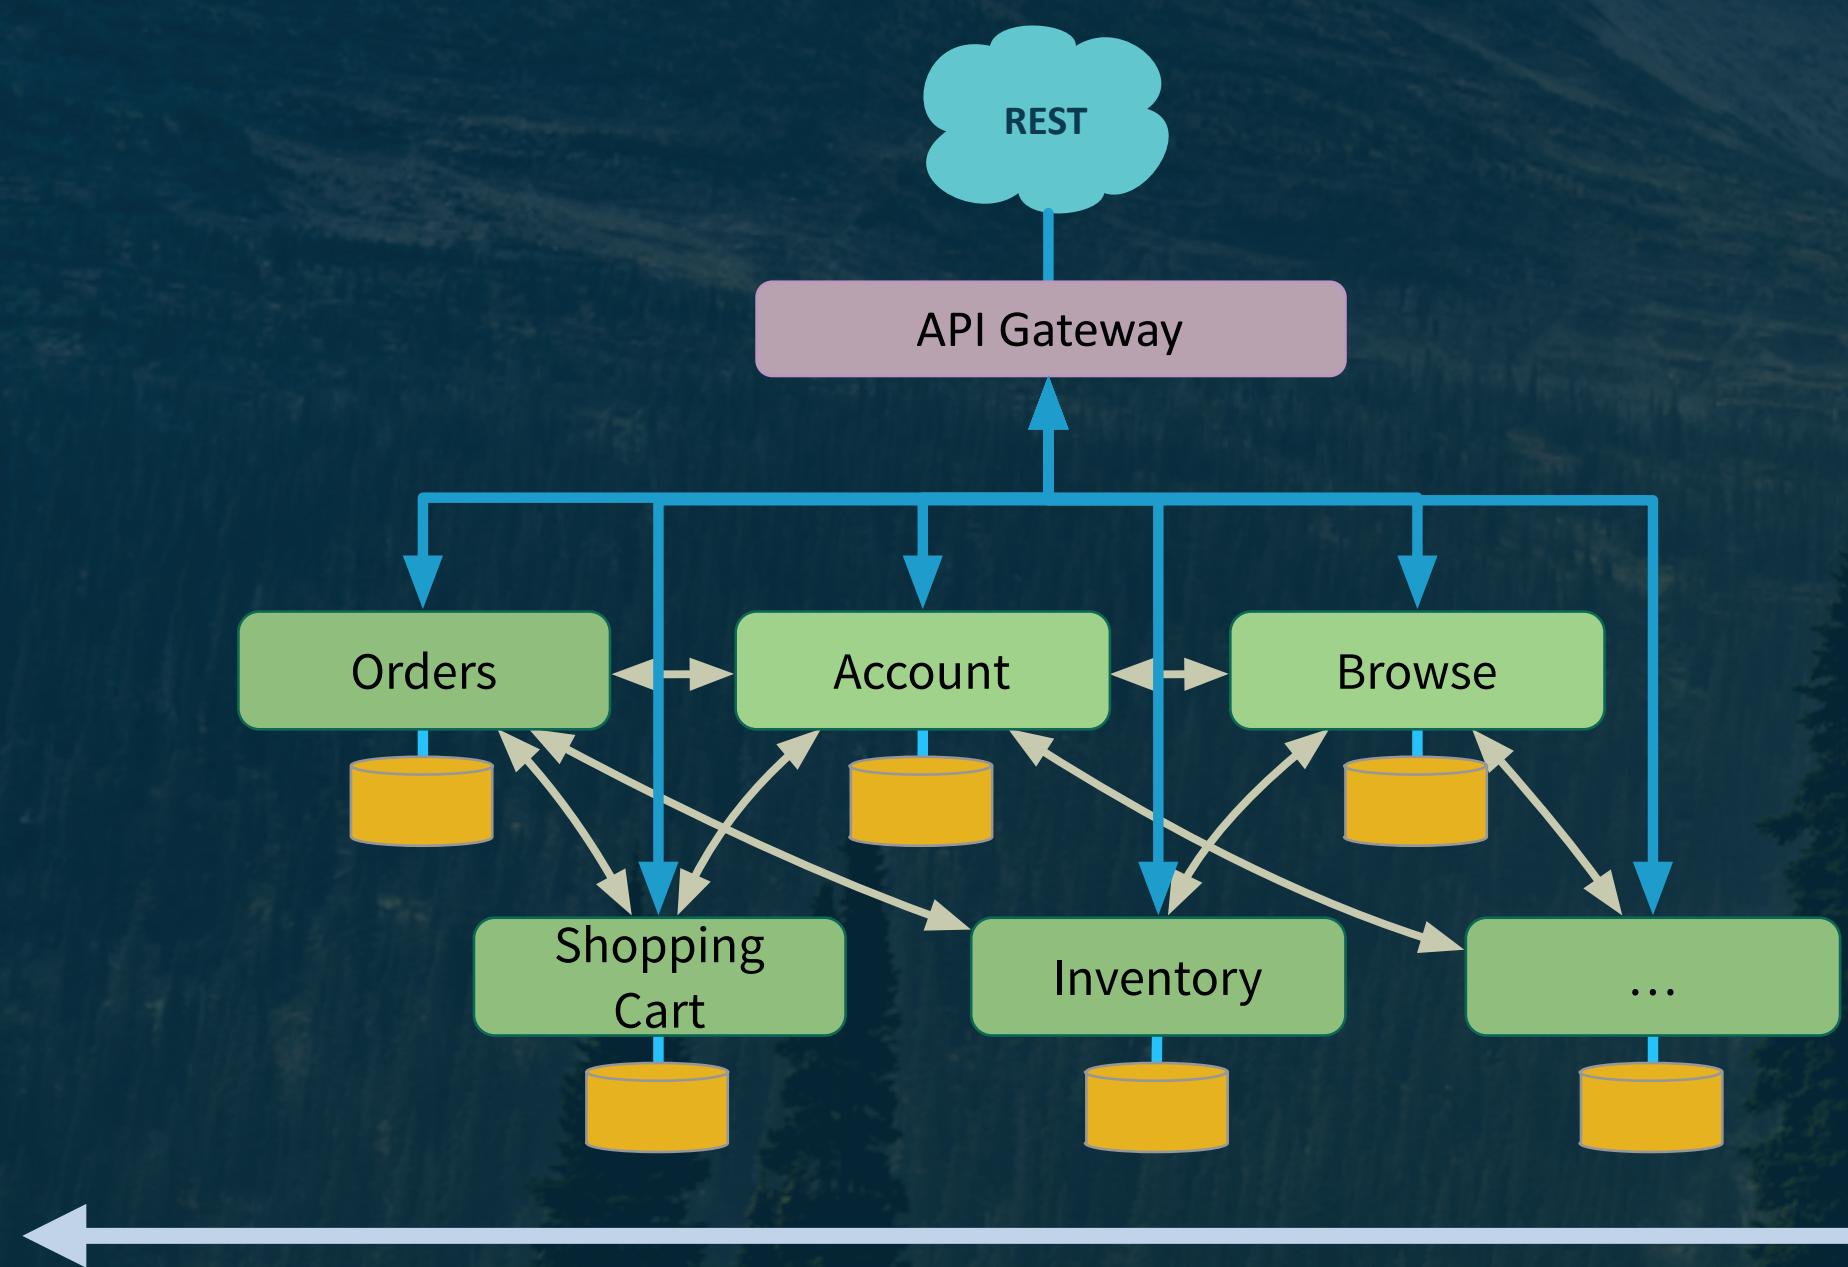
207

566

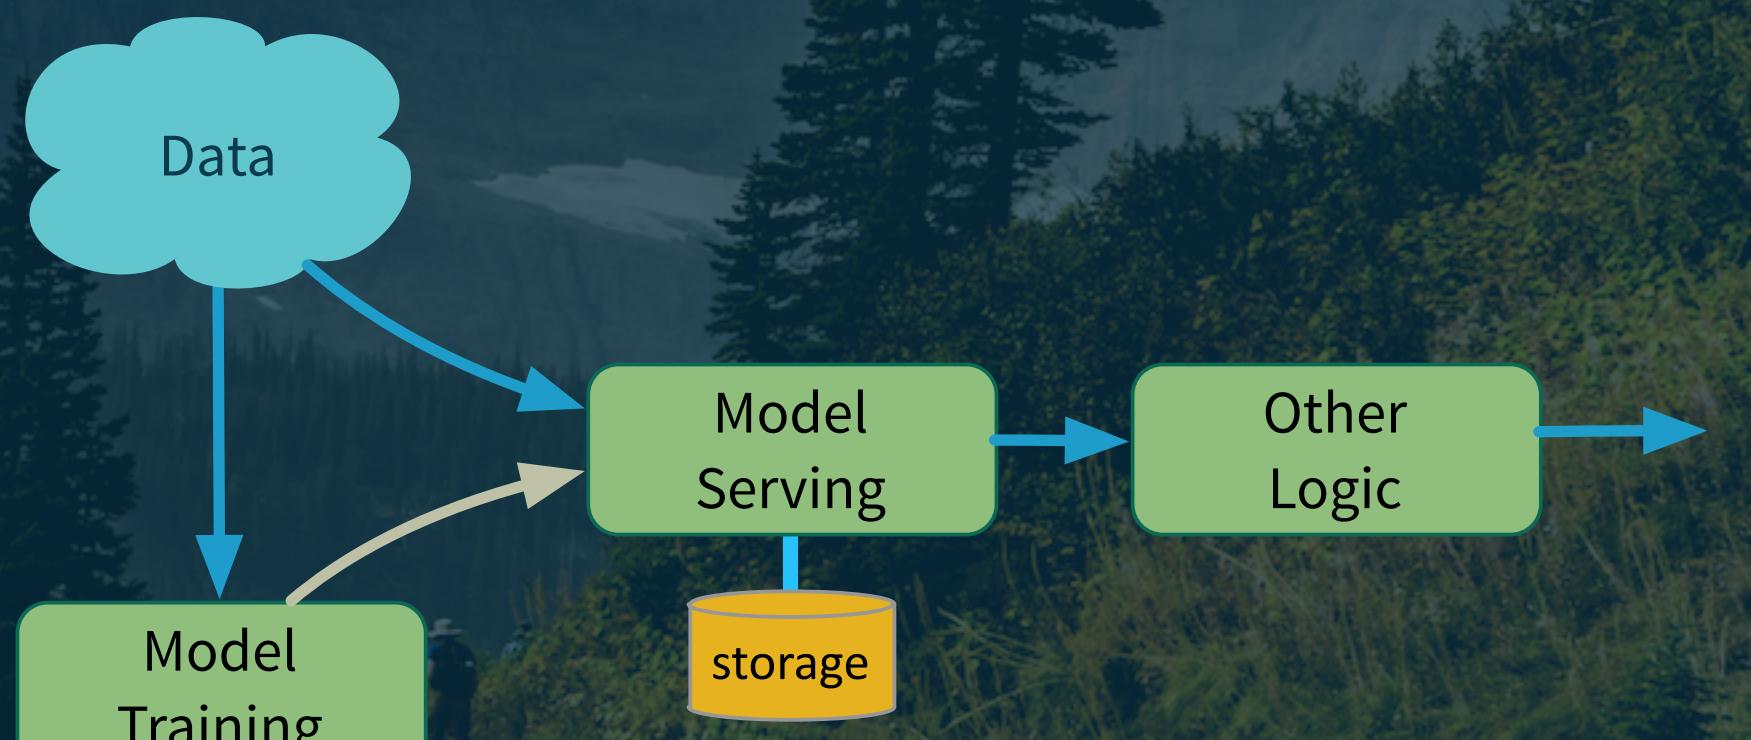


# A Spectrum of Microservices

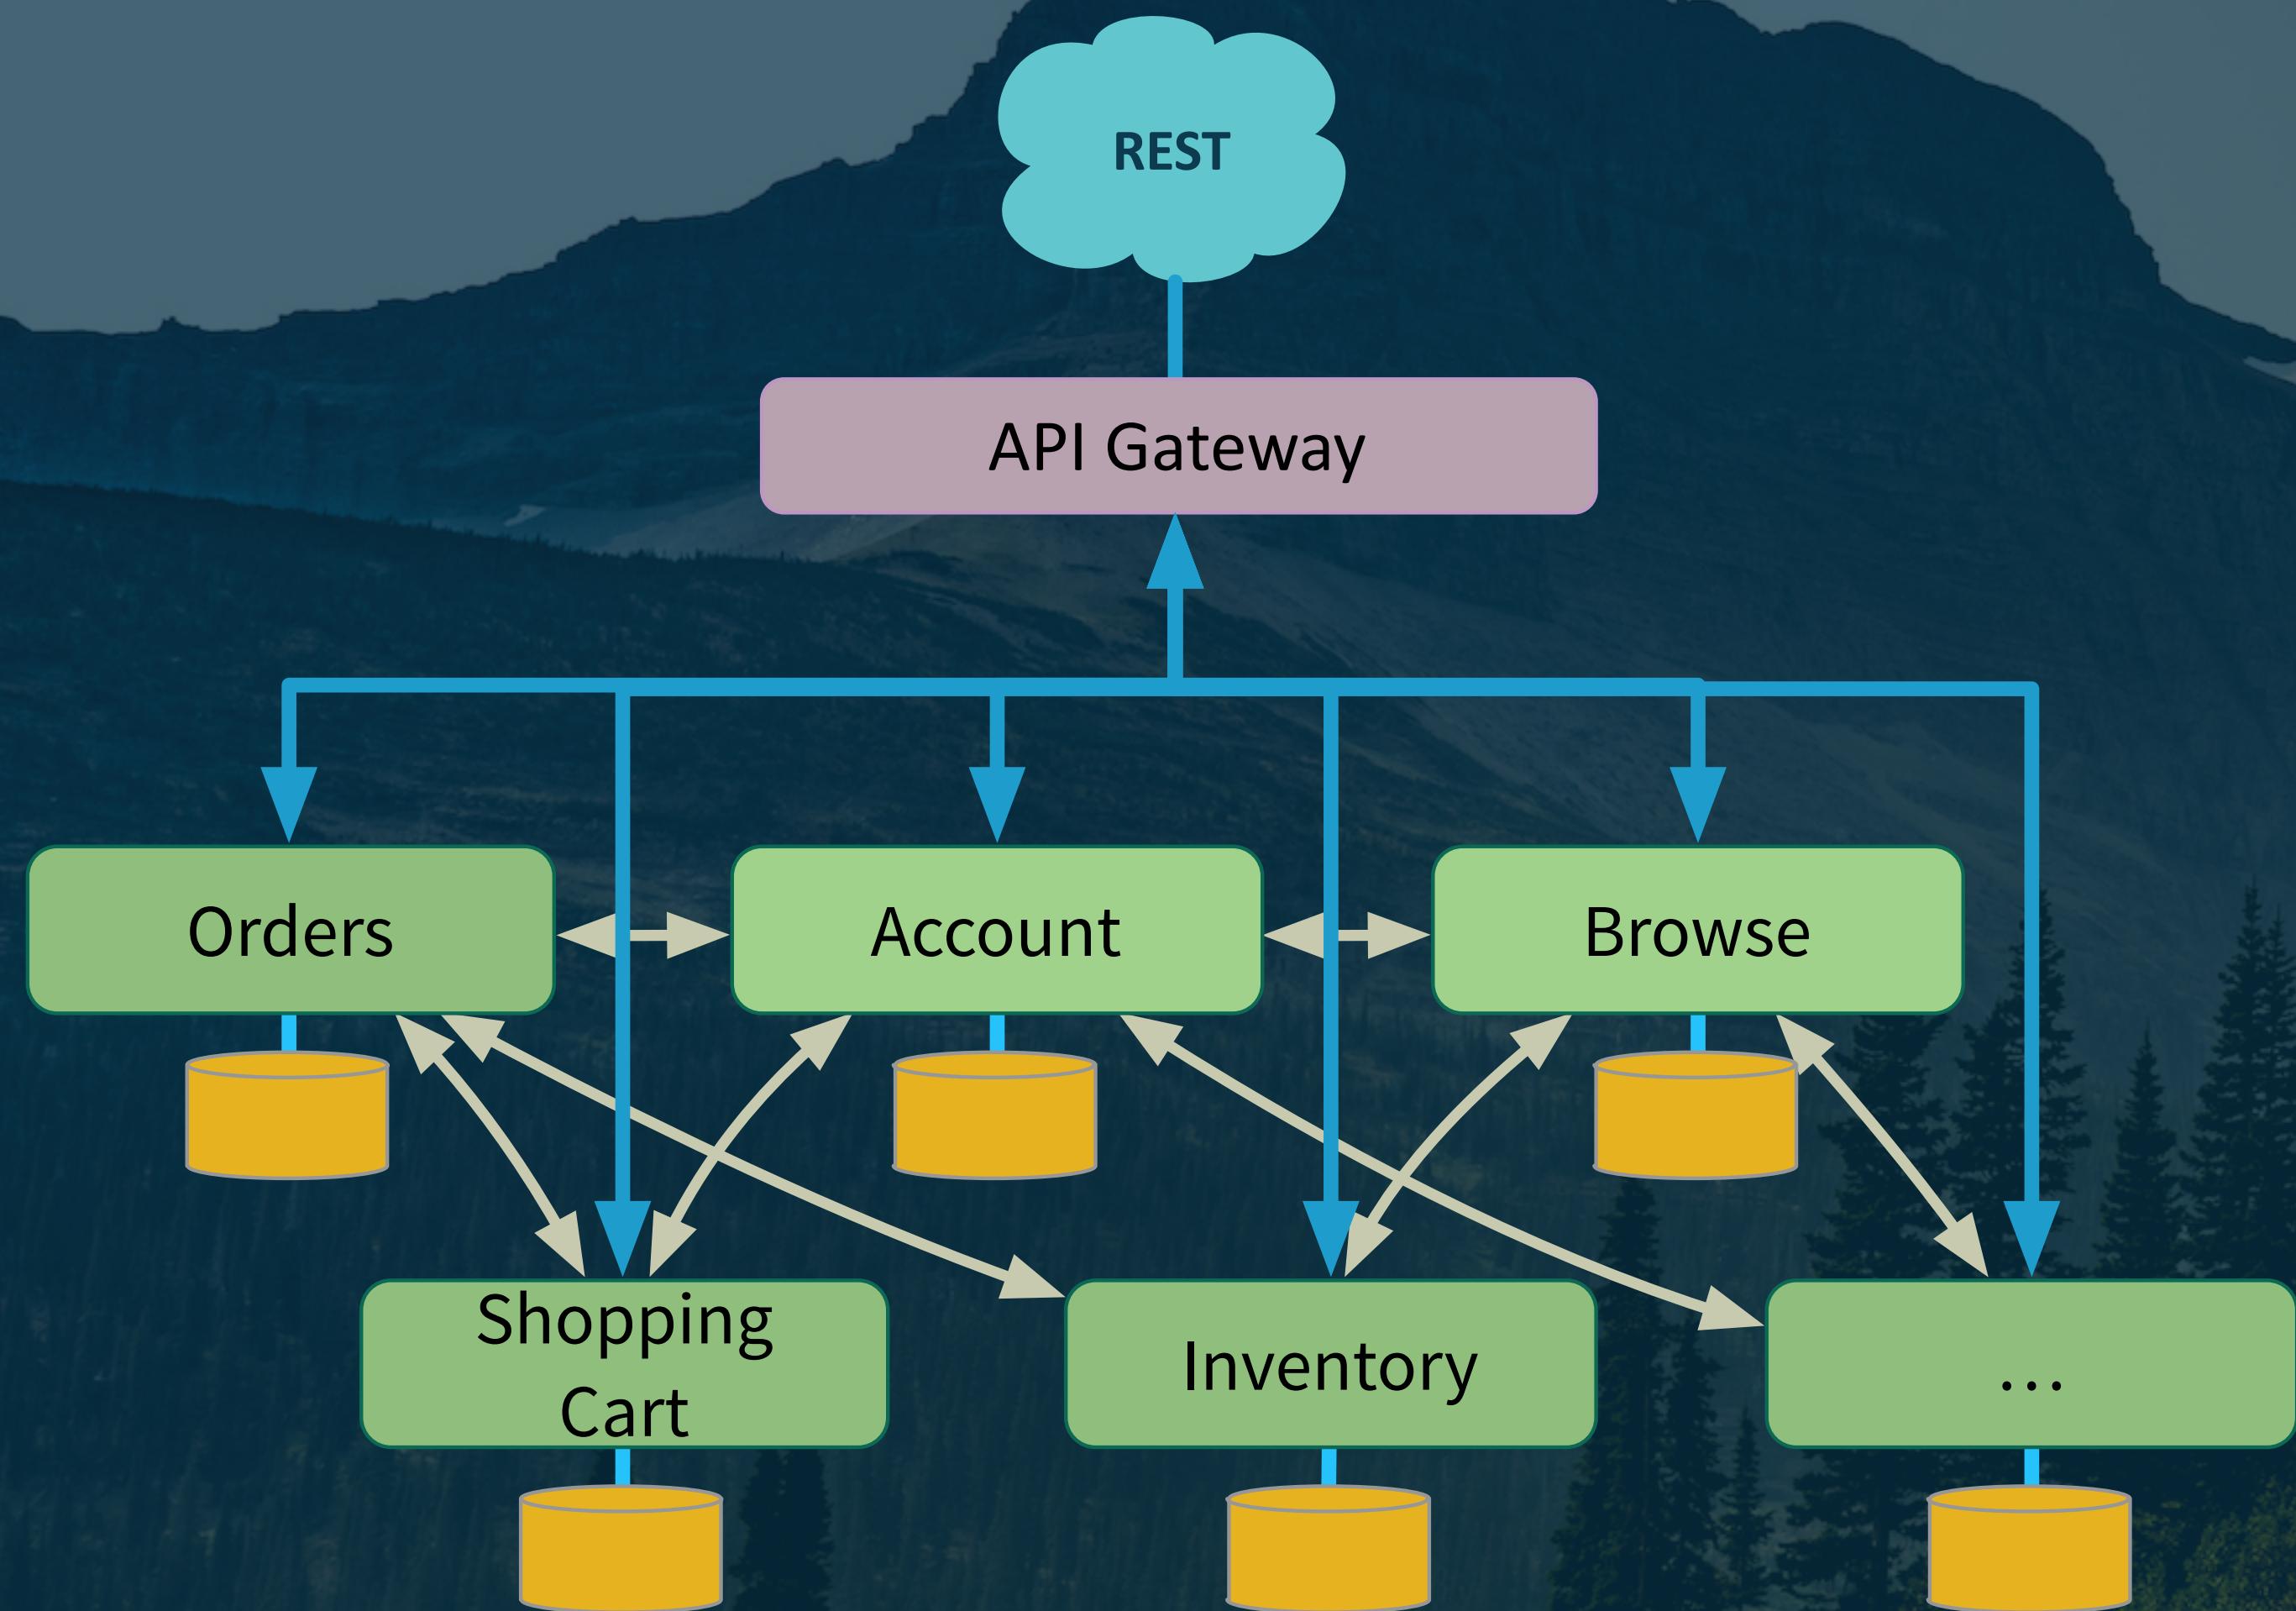
Event-driven μ-services



“Record-centric” μ-services



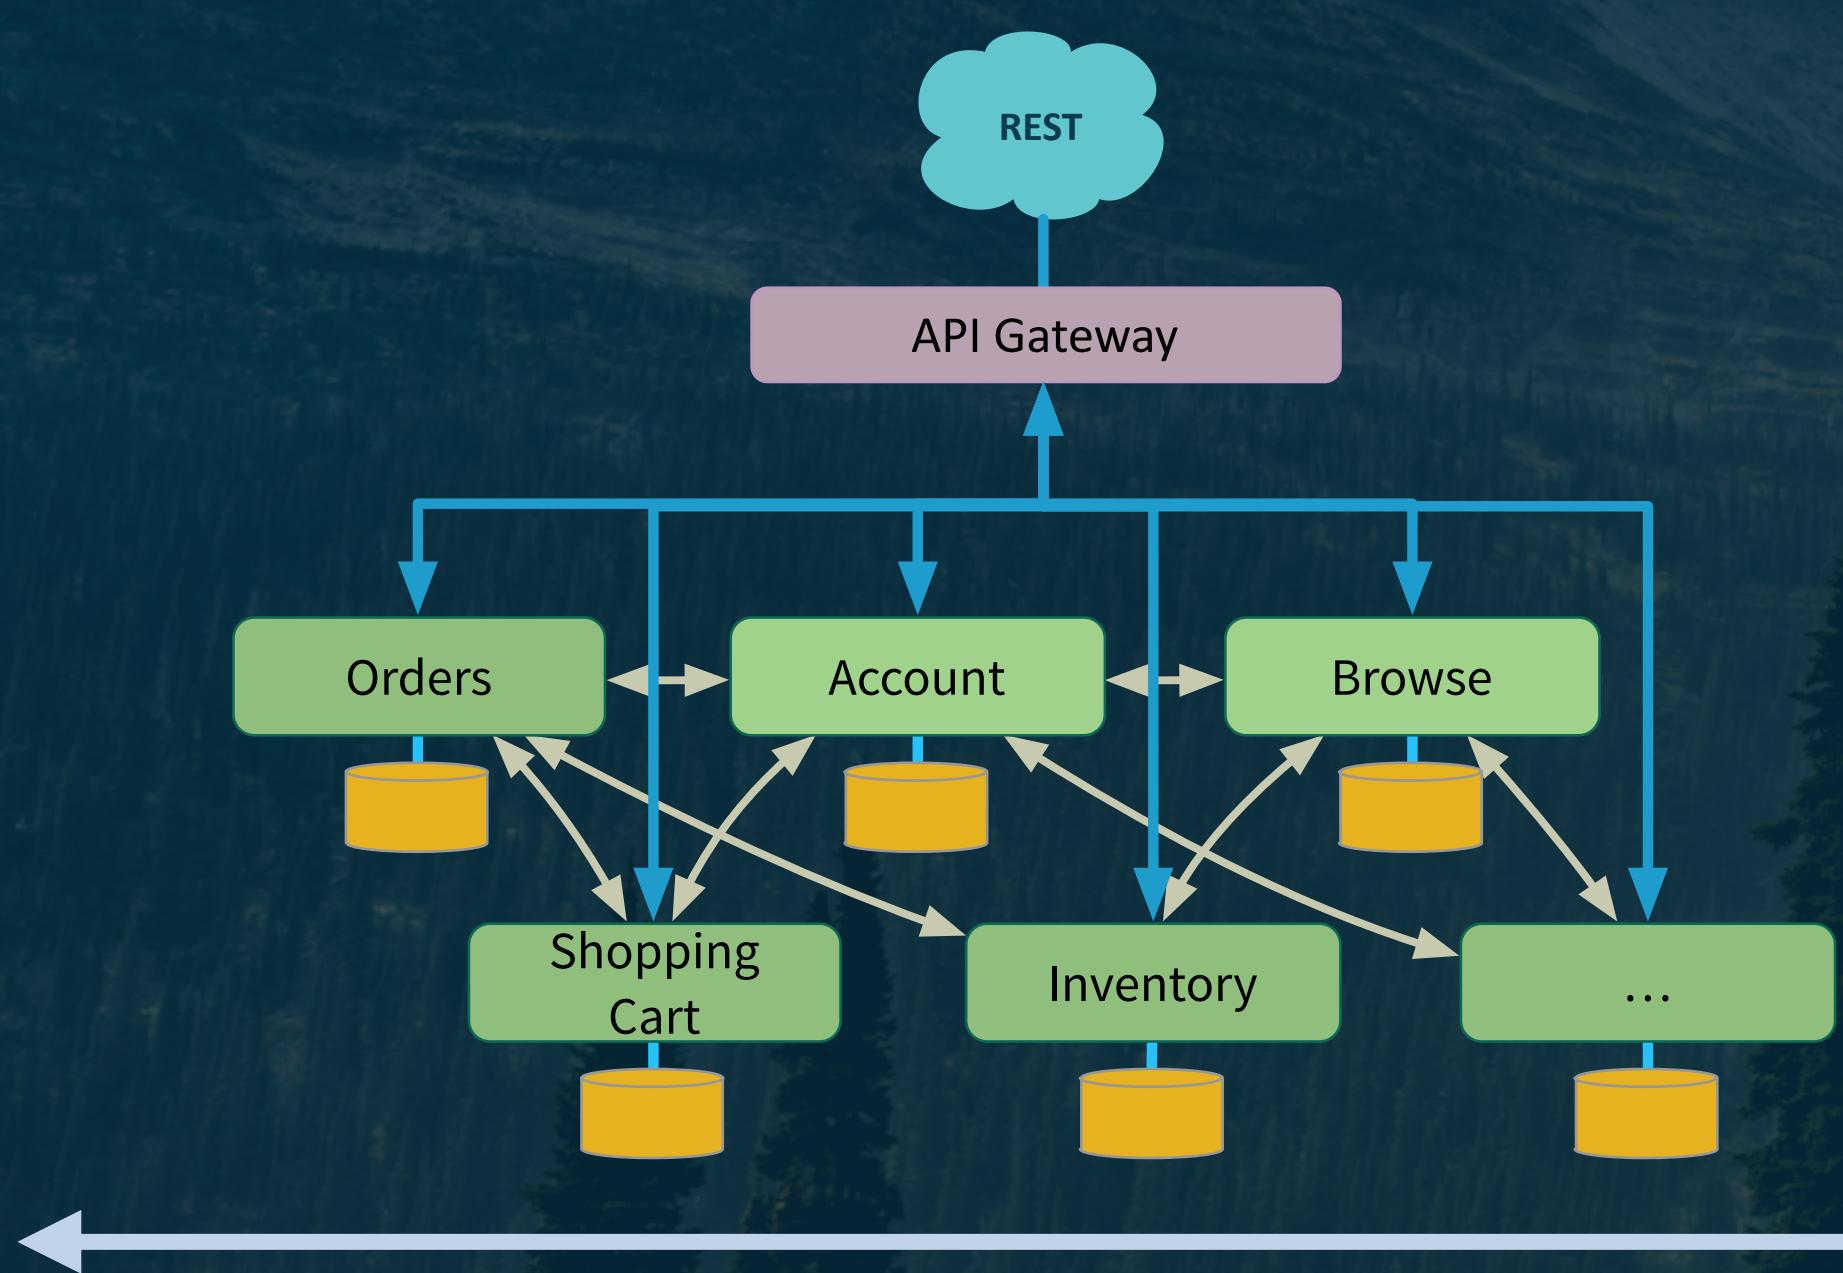
# A Spectrum of Microservices



- Each datum has an identity
- Process each one uniquely
- Think sessions and state machines

# A Spectrum of Microservices

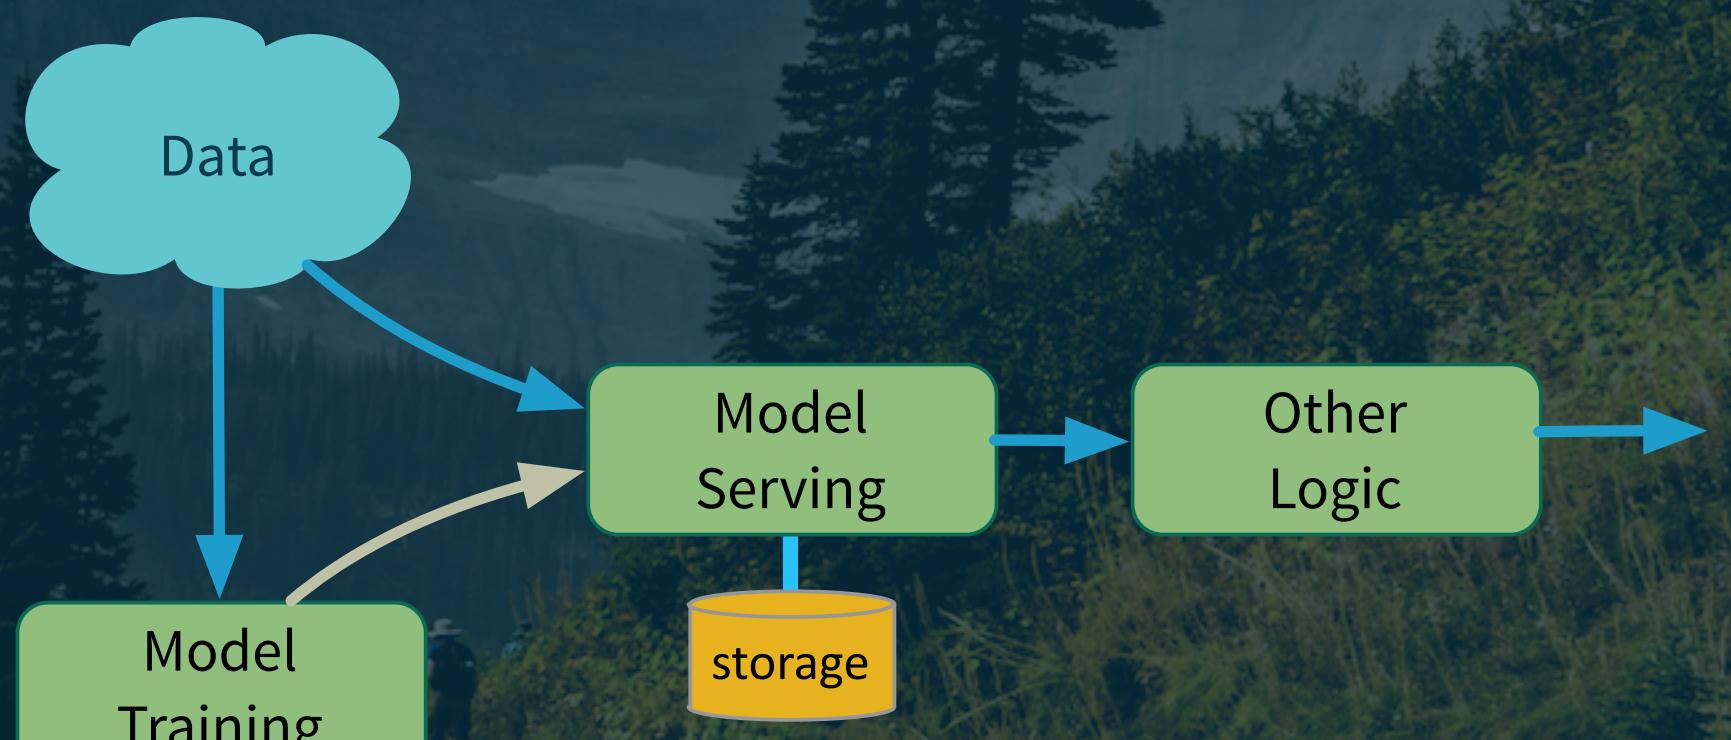
Event-driven μ-services



Events

@deanwampler

“Record-centric” μ-services

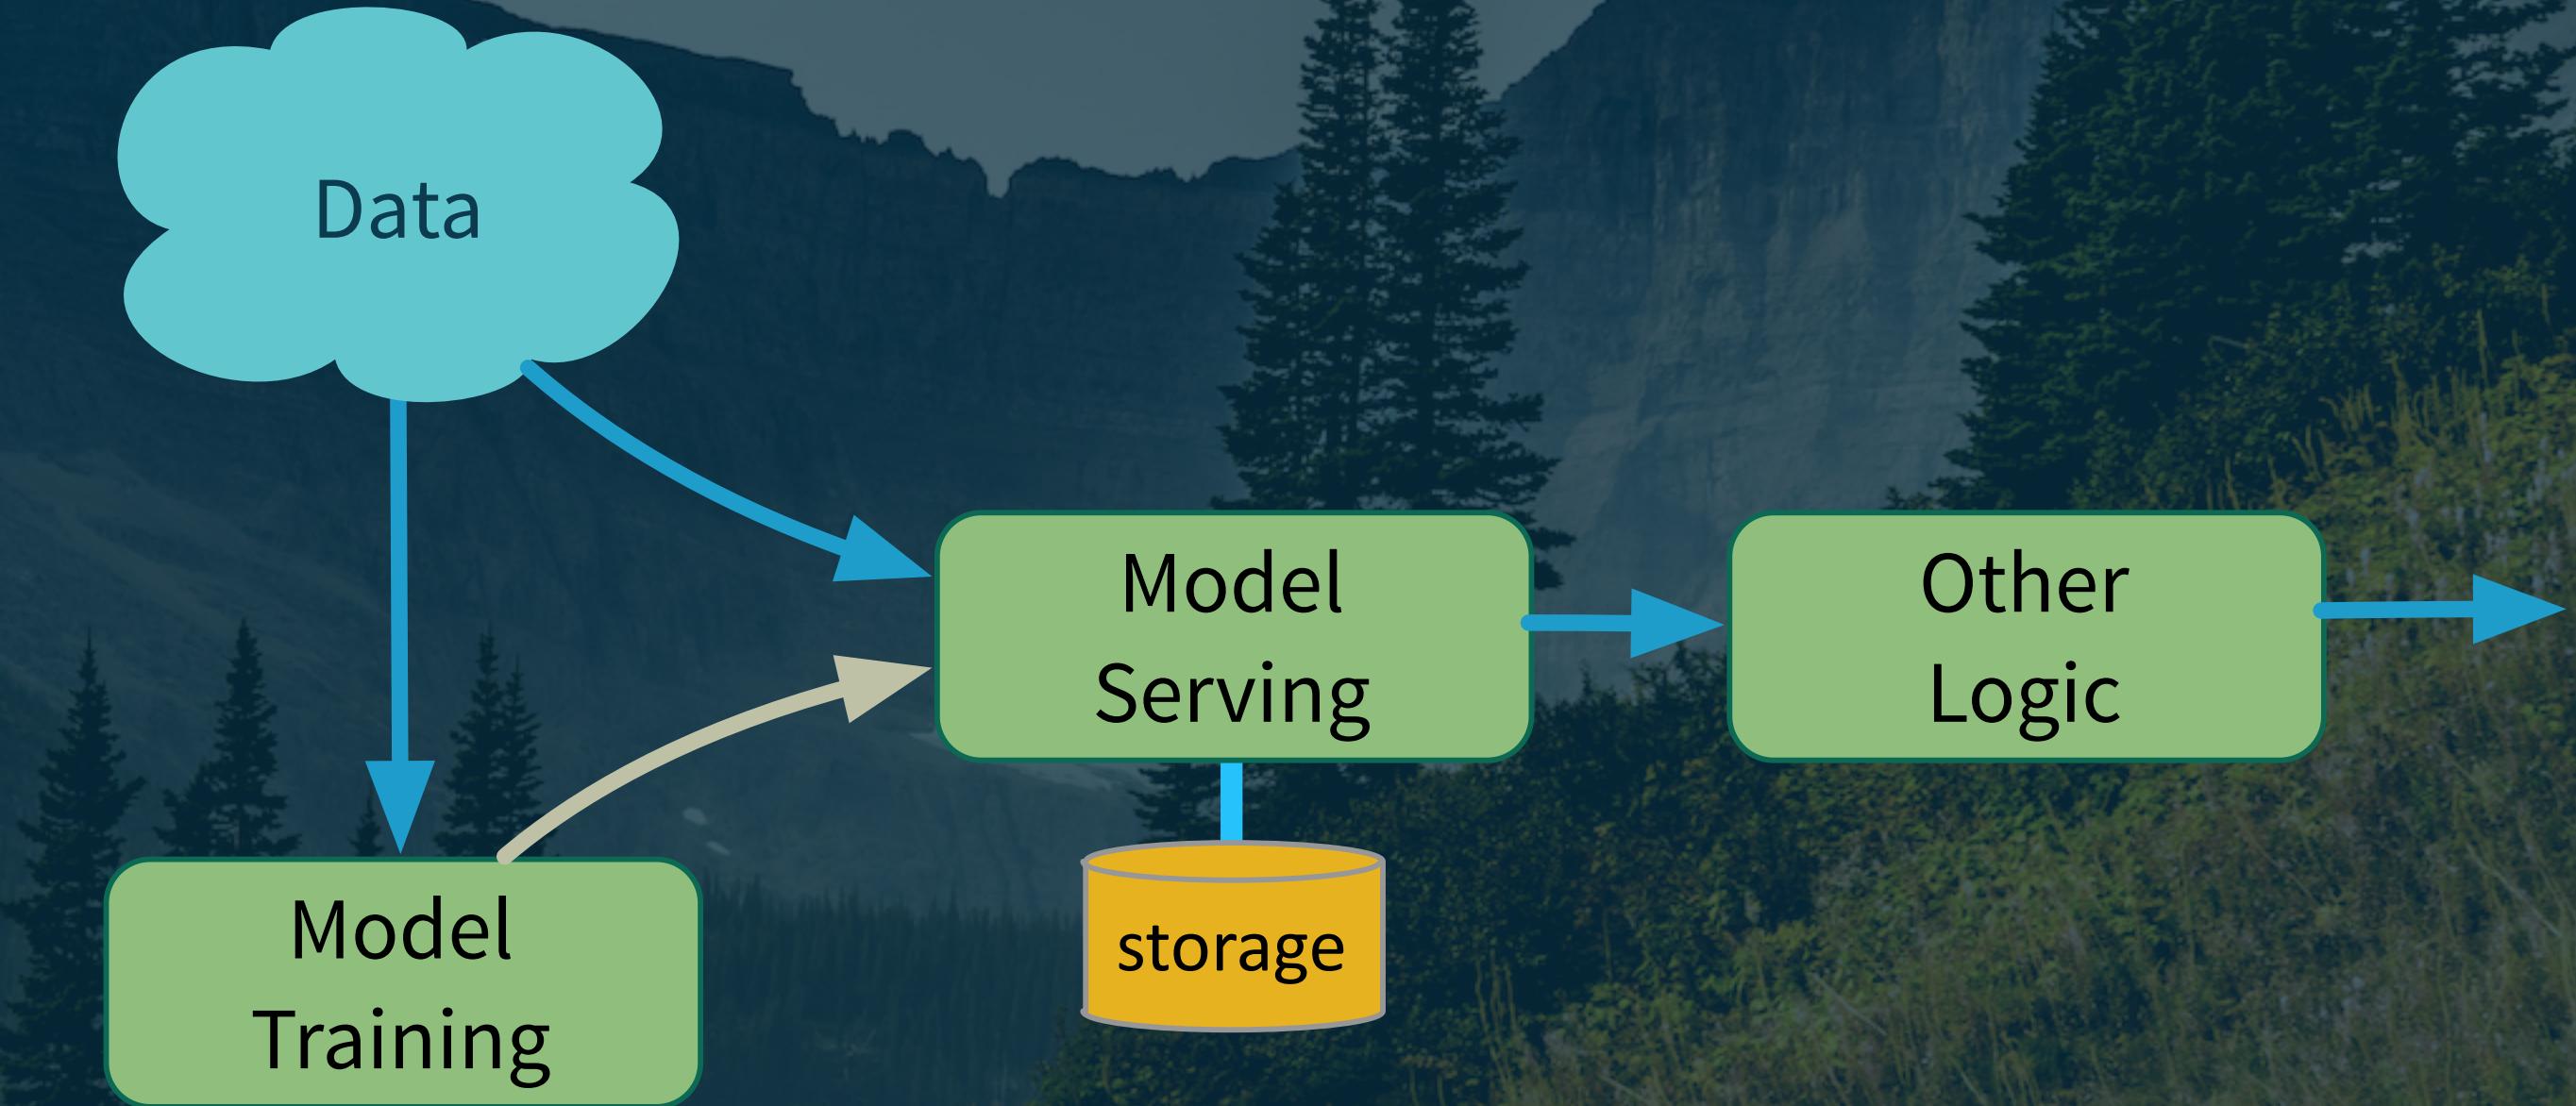


Records

← Data Spectrum →

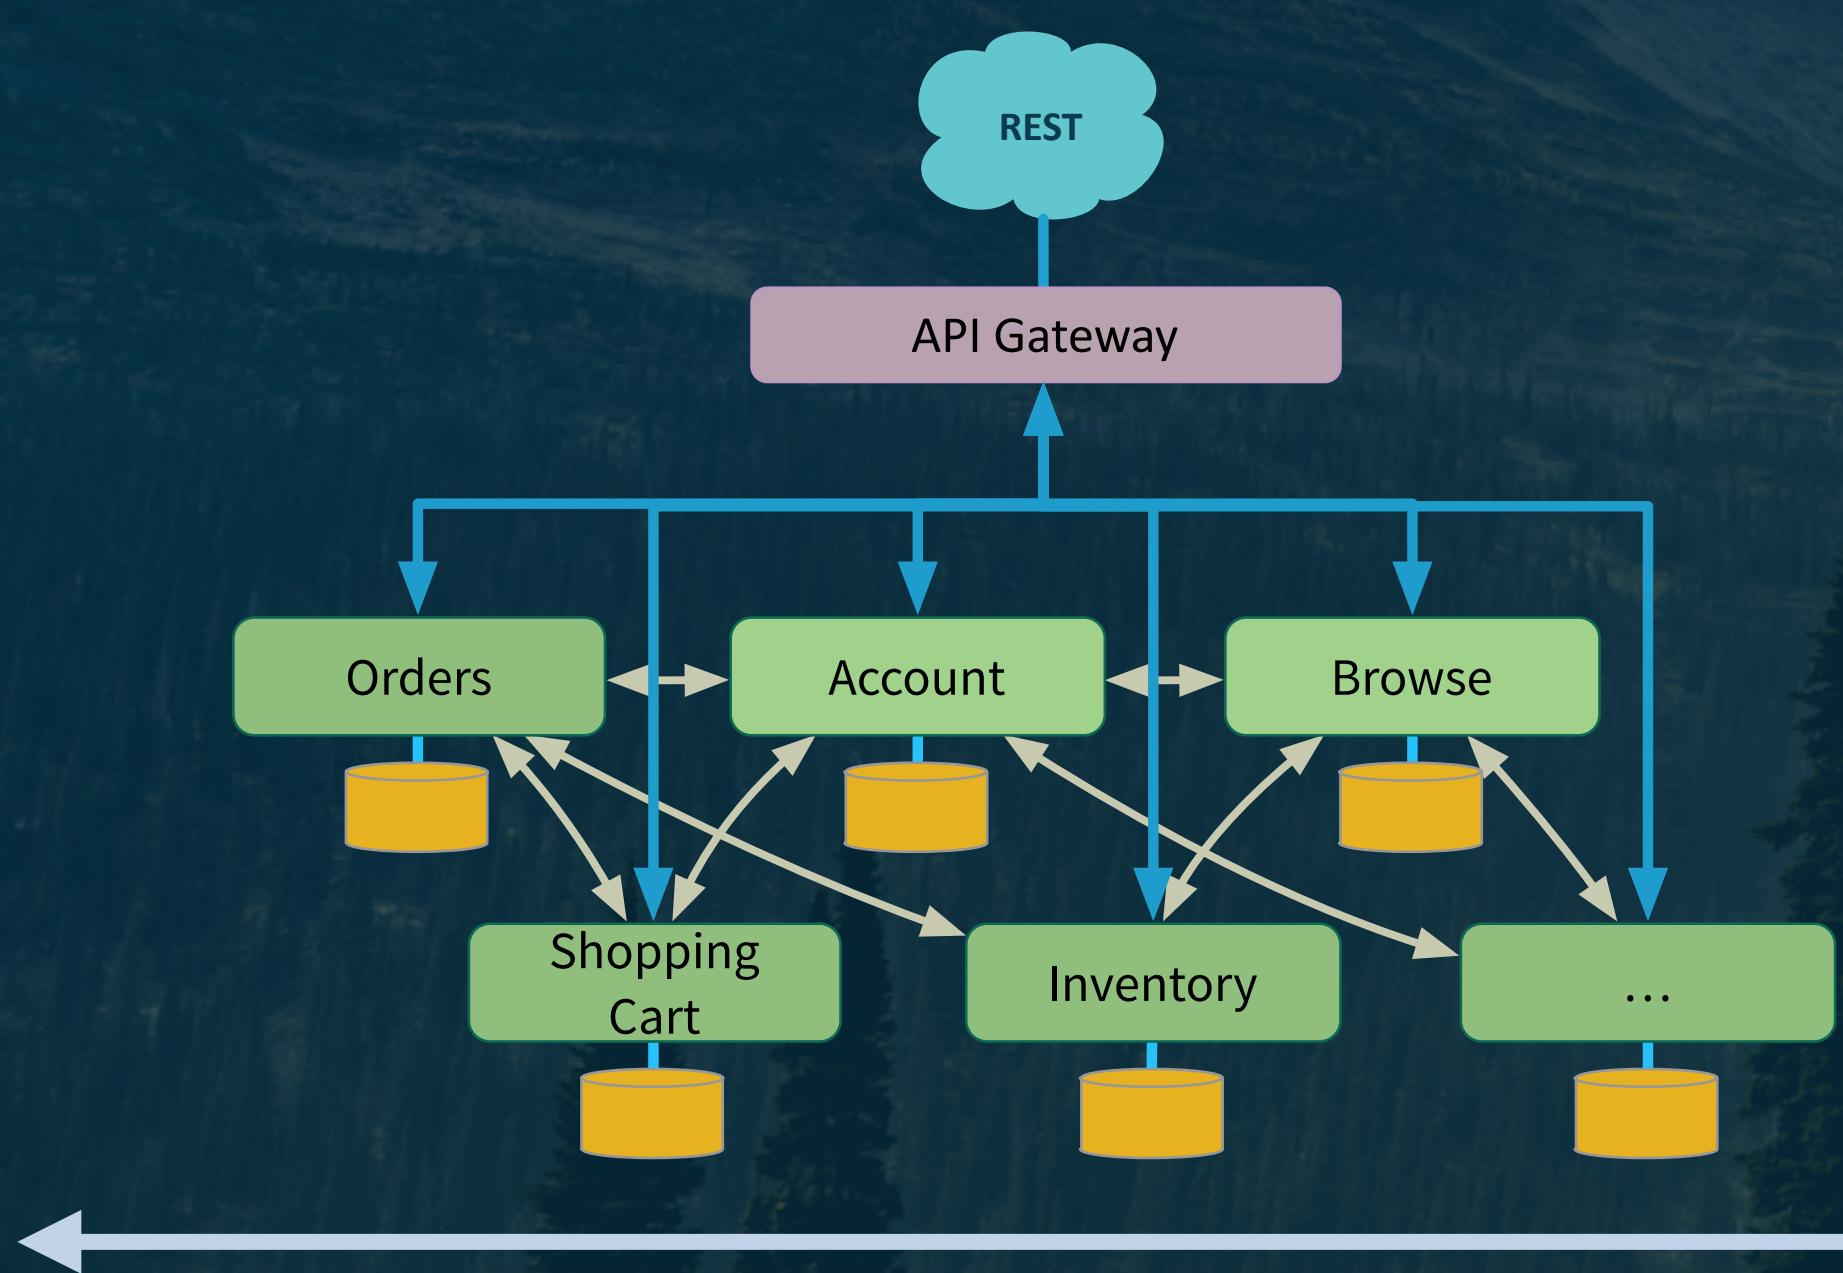
# A Spectrum of Microservices

- “Anonymous” records
- Process *en masse*
- Think SQL queries for analytics

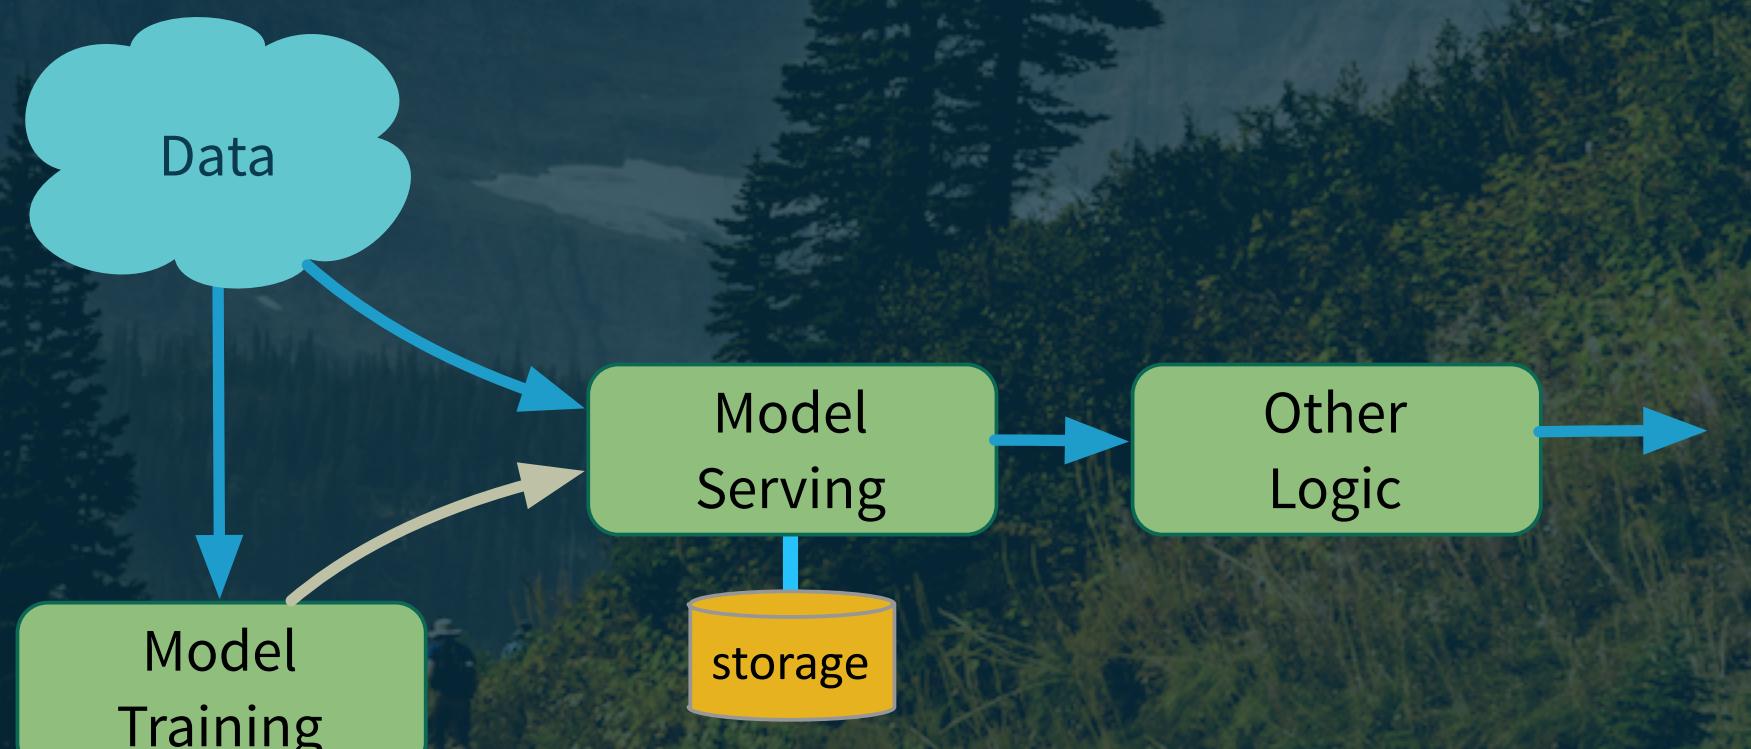


# A Spectrum of Microservices

Event-driven μ-services



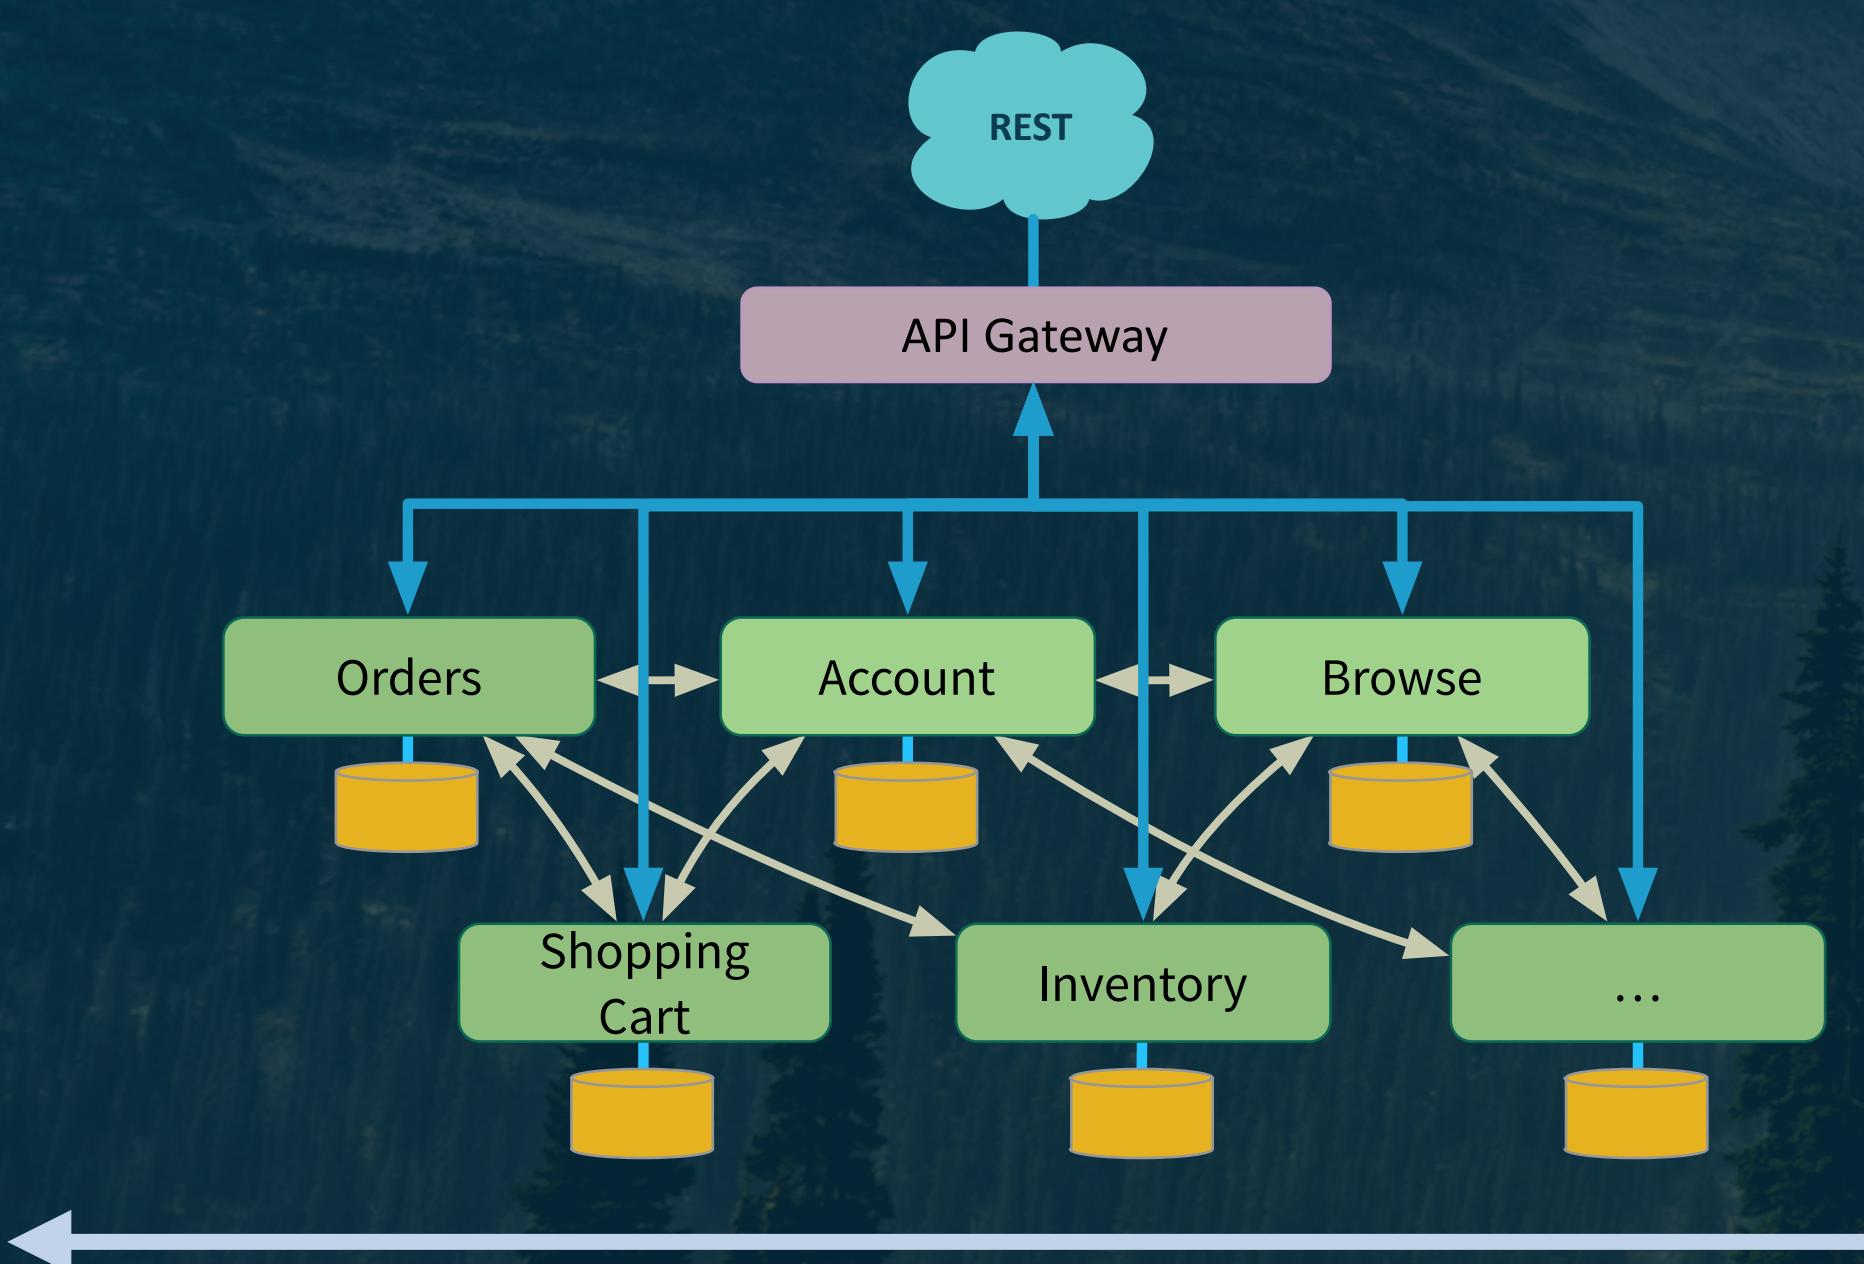
“Record-centric” μ-services



# A Spectrum of Microservices



## Event-driven μ-services



Events  
@deanwampler

Akka emerged from the left-hand side of the spectrum, the world of highly *Reactive* microservices.

Akka Streams pushes to the right, more data-centric.

<https://www.reactivemanifesto.org/>

← Data Spectrum →

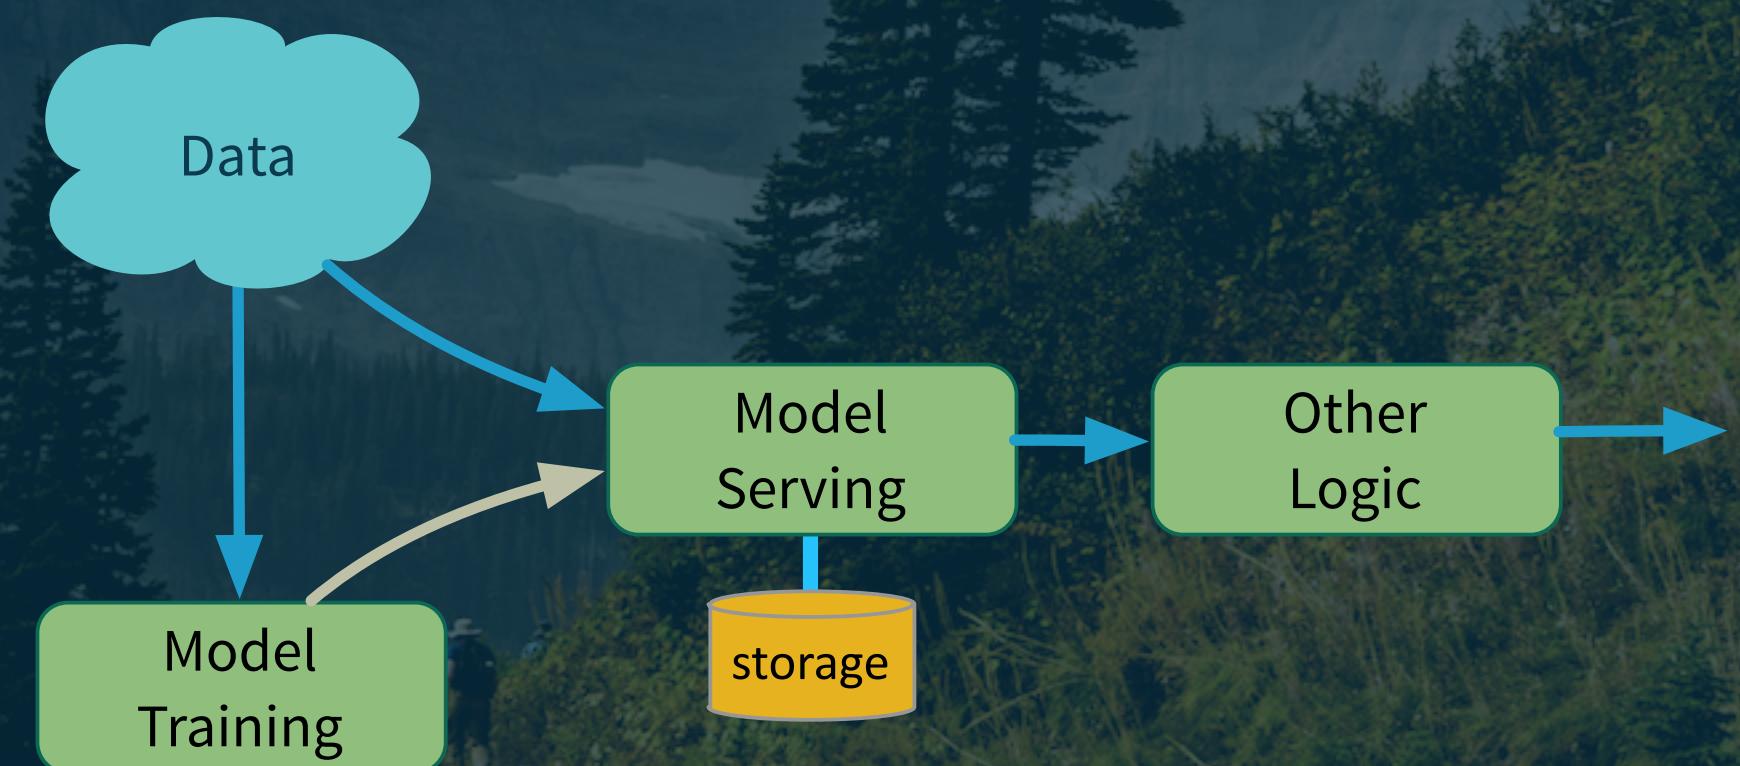
Records

# A Spectrum of Microservices



Emerged from the right-hand “Record-centric”  $\mu$ -services side.

Kafka Streams pushes to the left, supporting many event-processing scenarios.



$\leftarrow$  Events  $\rightarrow$  Records  
 $\leftarrow$  Data Spectrum  $\rightarrow$

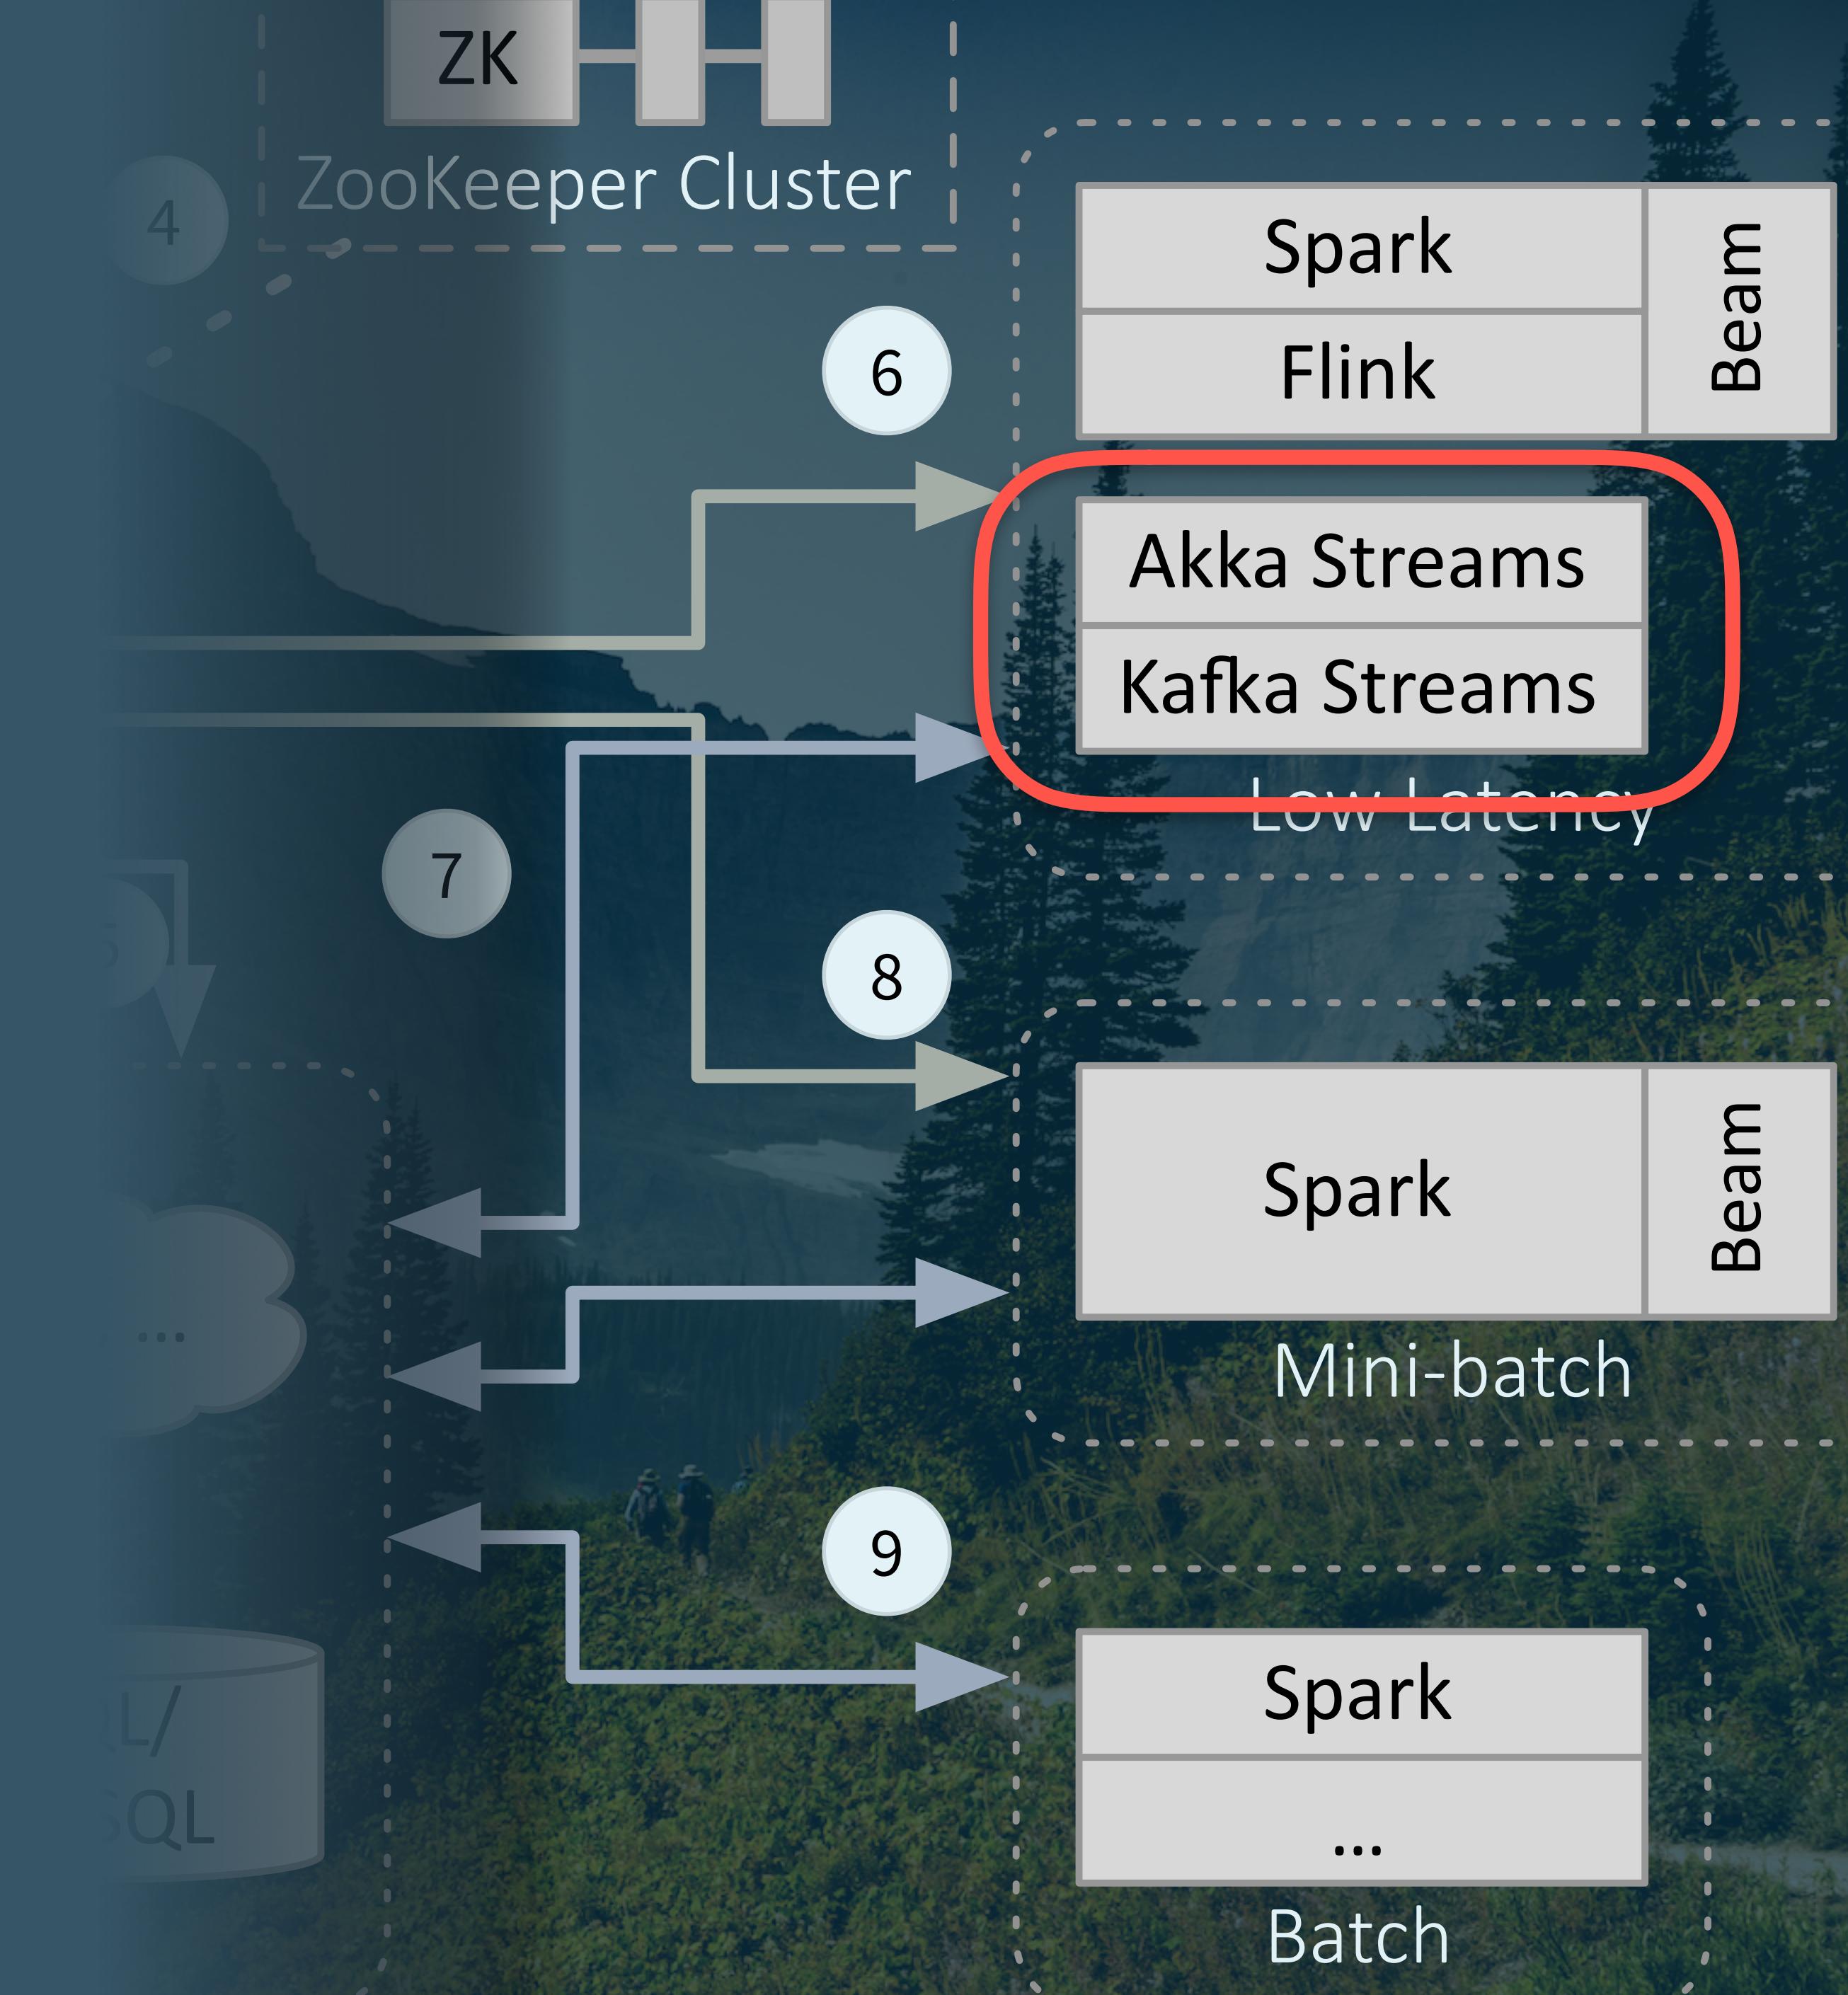
# Akka/Kafka Streams

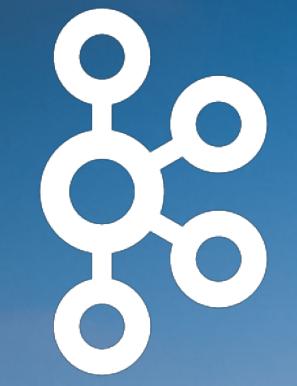
Wired:

- You already know microservices
- Low latency
- Flexibility

Tired:

- Automatic scalability
- Less-rich semantics





kafka

# Kafka Streams





kafka

# Kafka Streams

- Important stream-processing semantics:
  - Problem: The stream never ends
  - Problem: I need compute  $X$  per *unit time*  $Y$
  - Problem: I need compute  $X$  per *session*

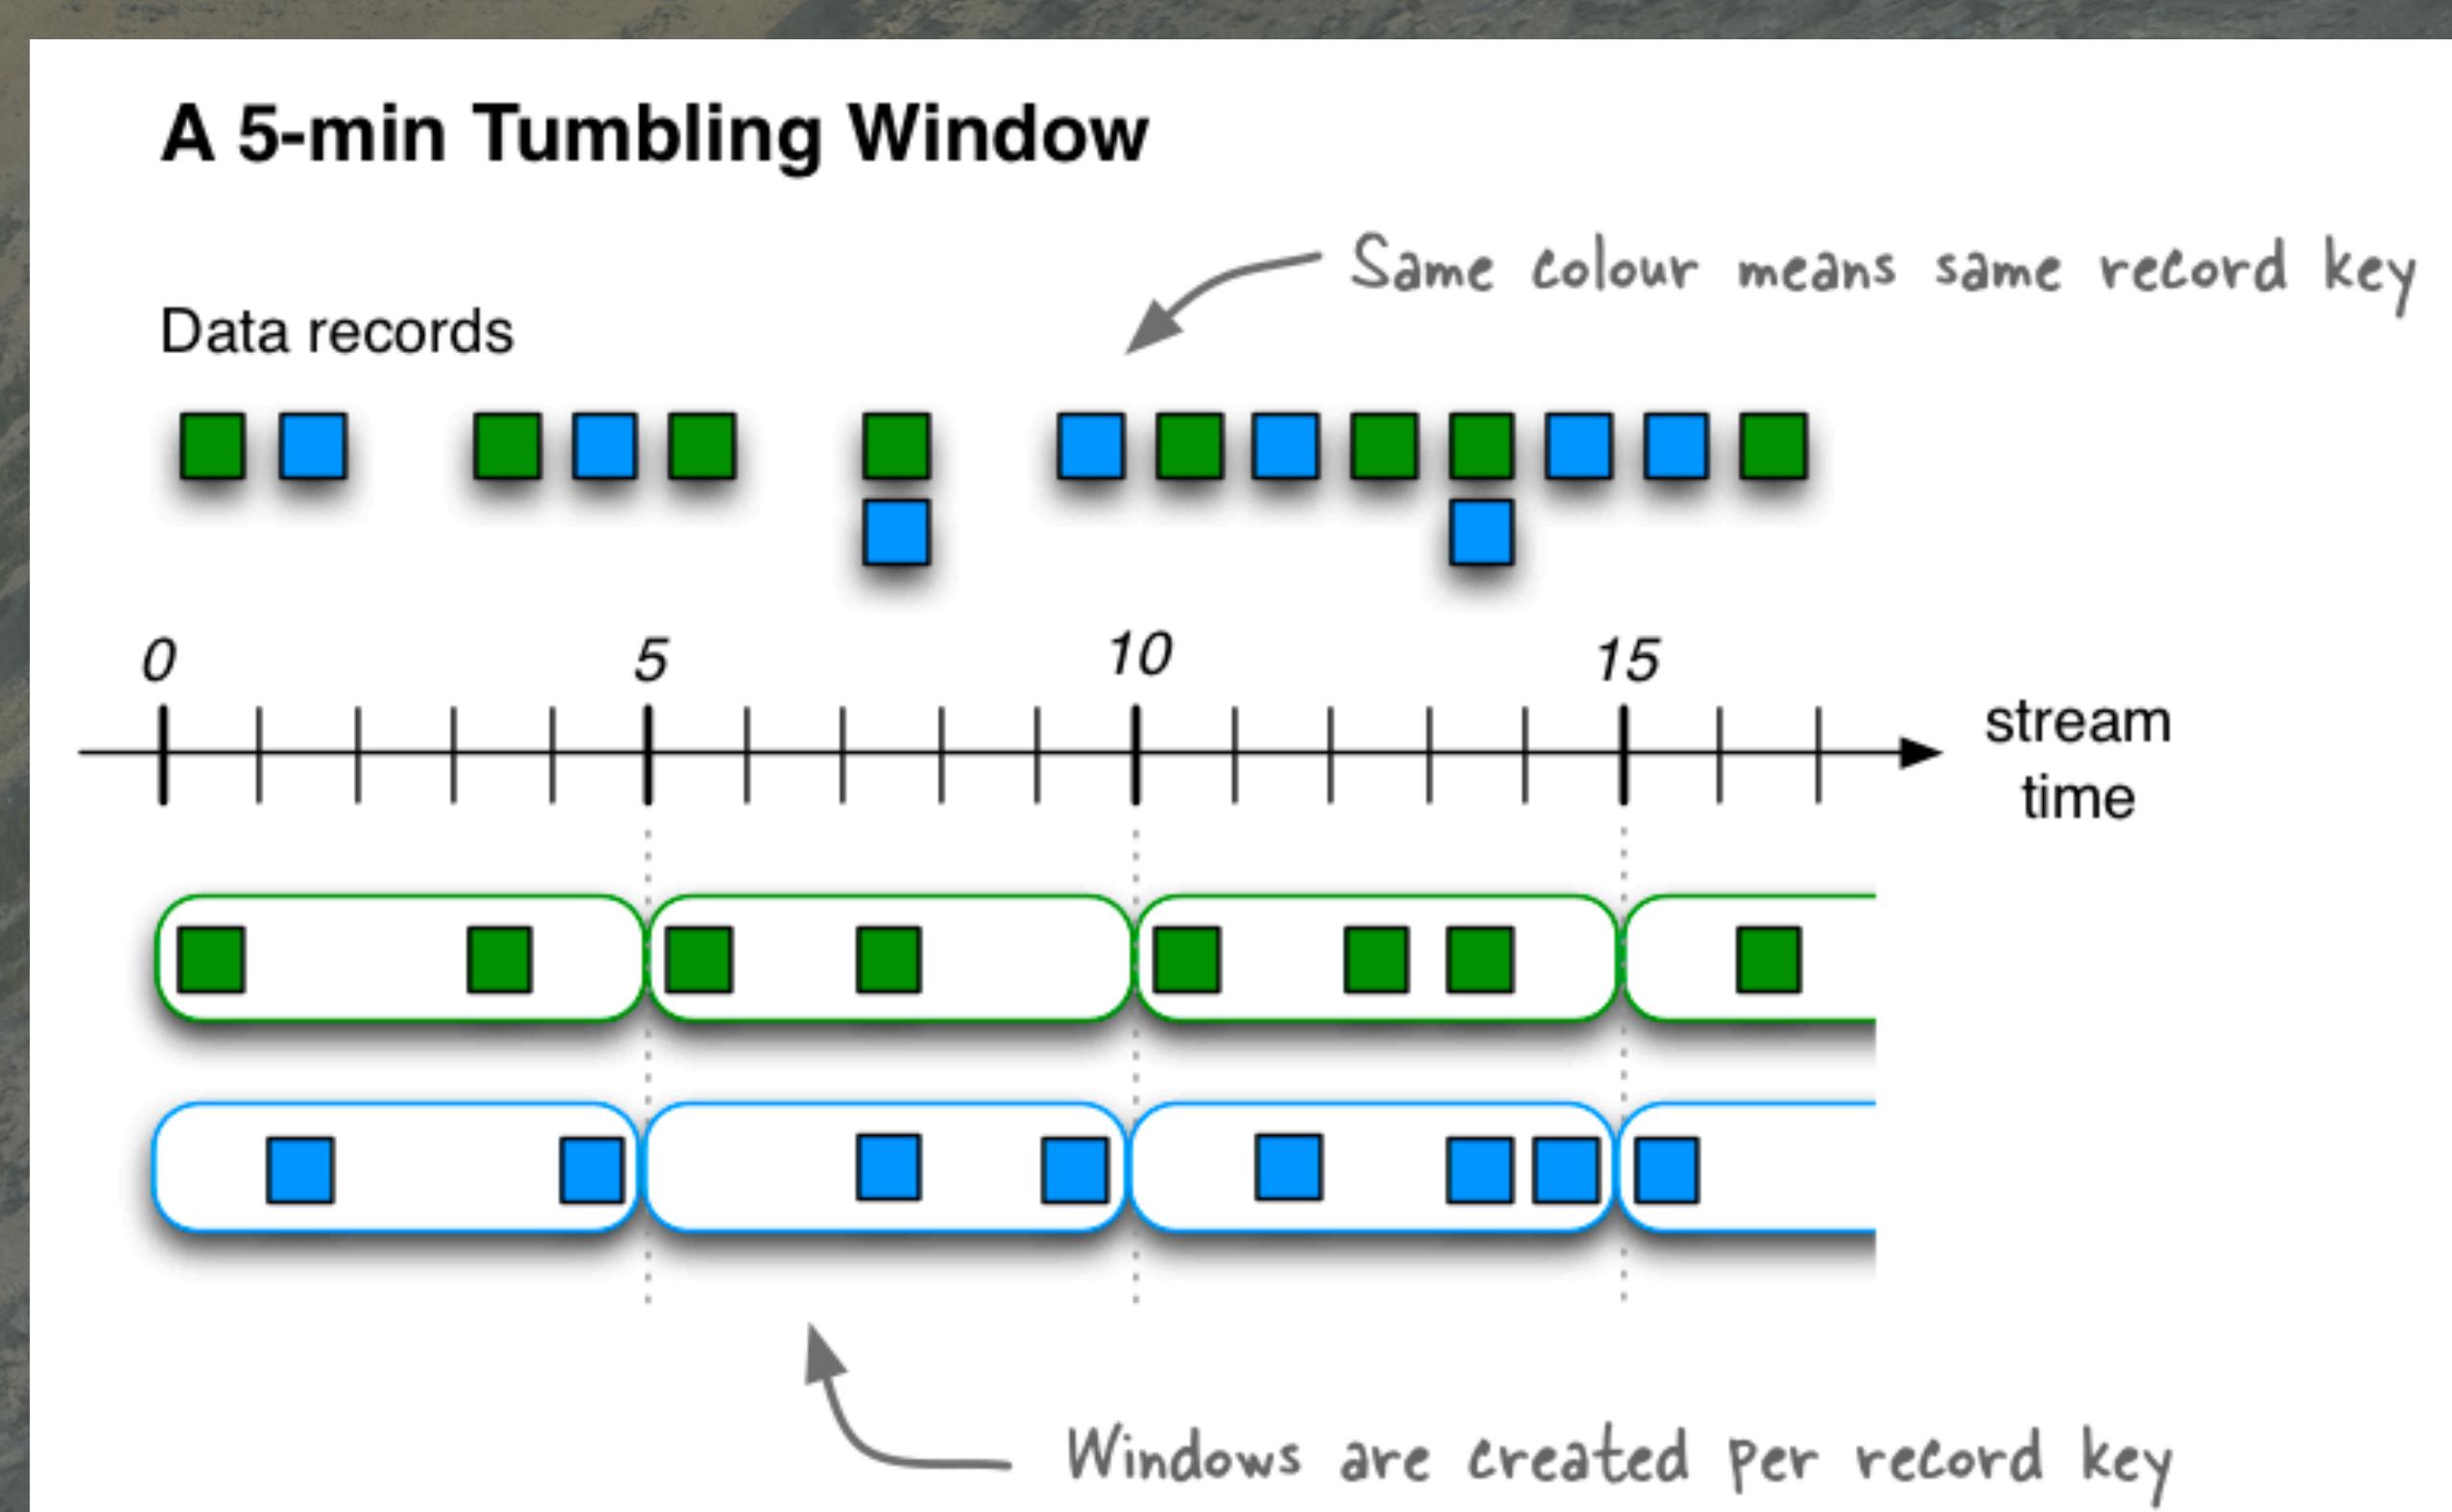


# kafka

# Kafka Streams

- A scenic view of mountains under a cloudy sky, serving as a background for a presentation slide.
- Important stream-processing semantics:
  - Solution: Windowing
  - Tumbling windows:

See my O'Reilly  
report for details.





# Kafka Streams

kafka

- Important stream-processing semantics:
- Solution: Windowing
  - Session windows:

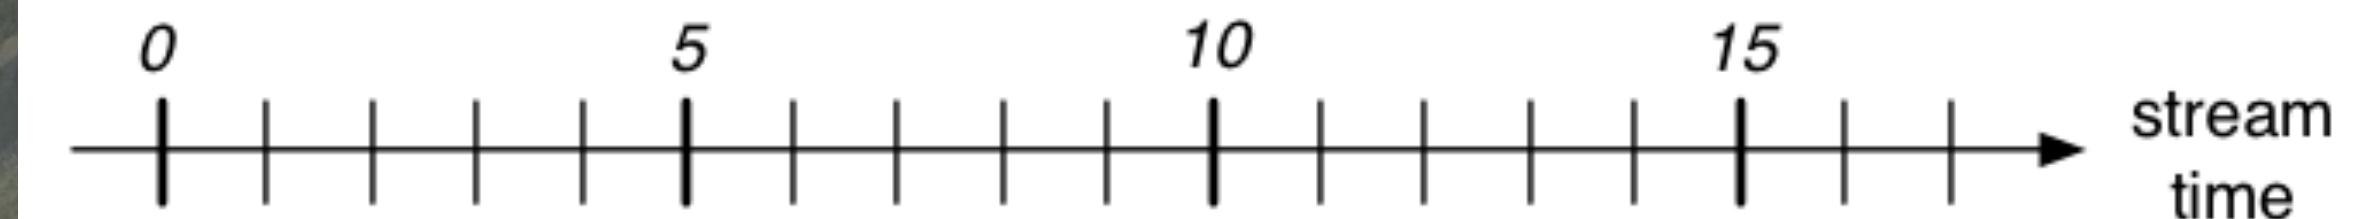
See my O'Reilly  
report for details.

## A Session Window with a 5-min inactivity gap

Data records



Same colour means same record key



Windows are created per record key



kafka

# Kafka Streams

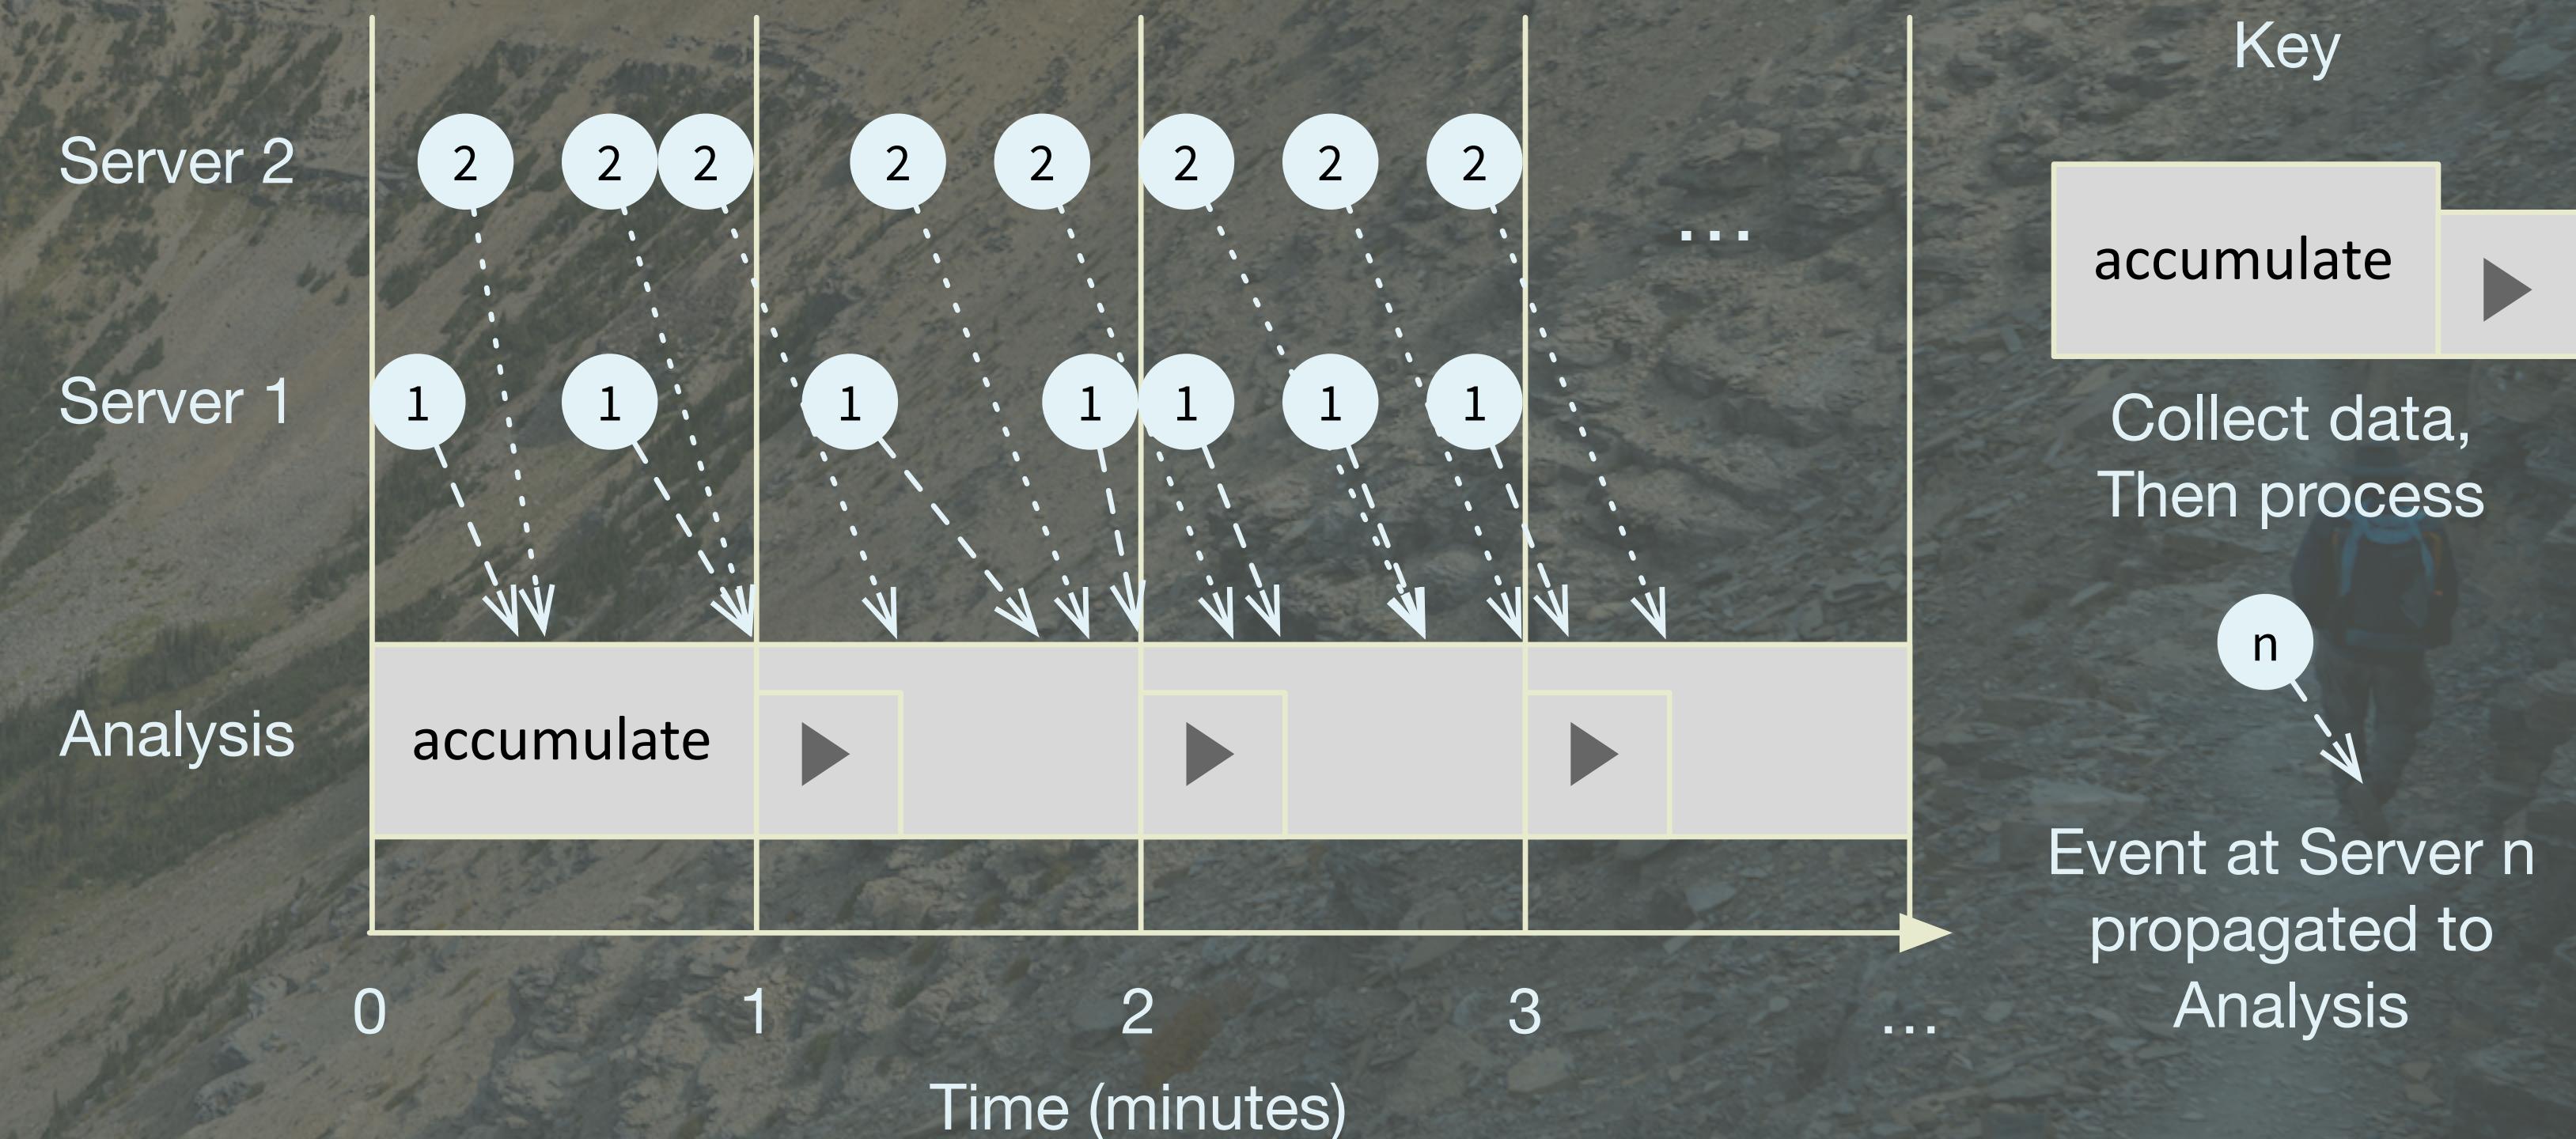
- Important stream-processing semantics:
  - Problem: When I say ... *per unit time Y*, I mean *event time*



kafka

# Kafka Streams

- Important stream-processing semantics:
- Solution: *event time* and *processing time*





# Kafka Streams

kafka

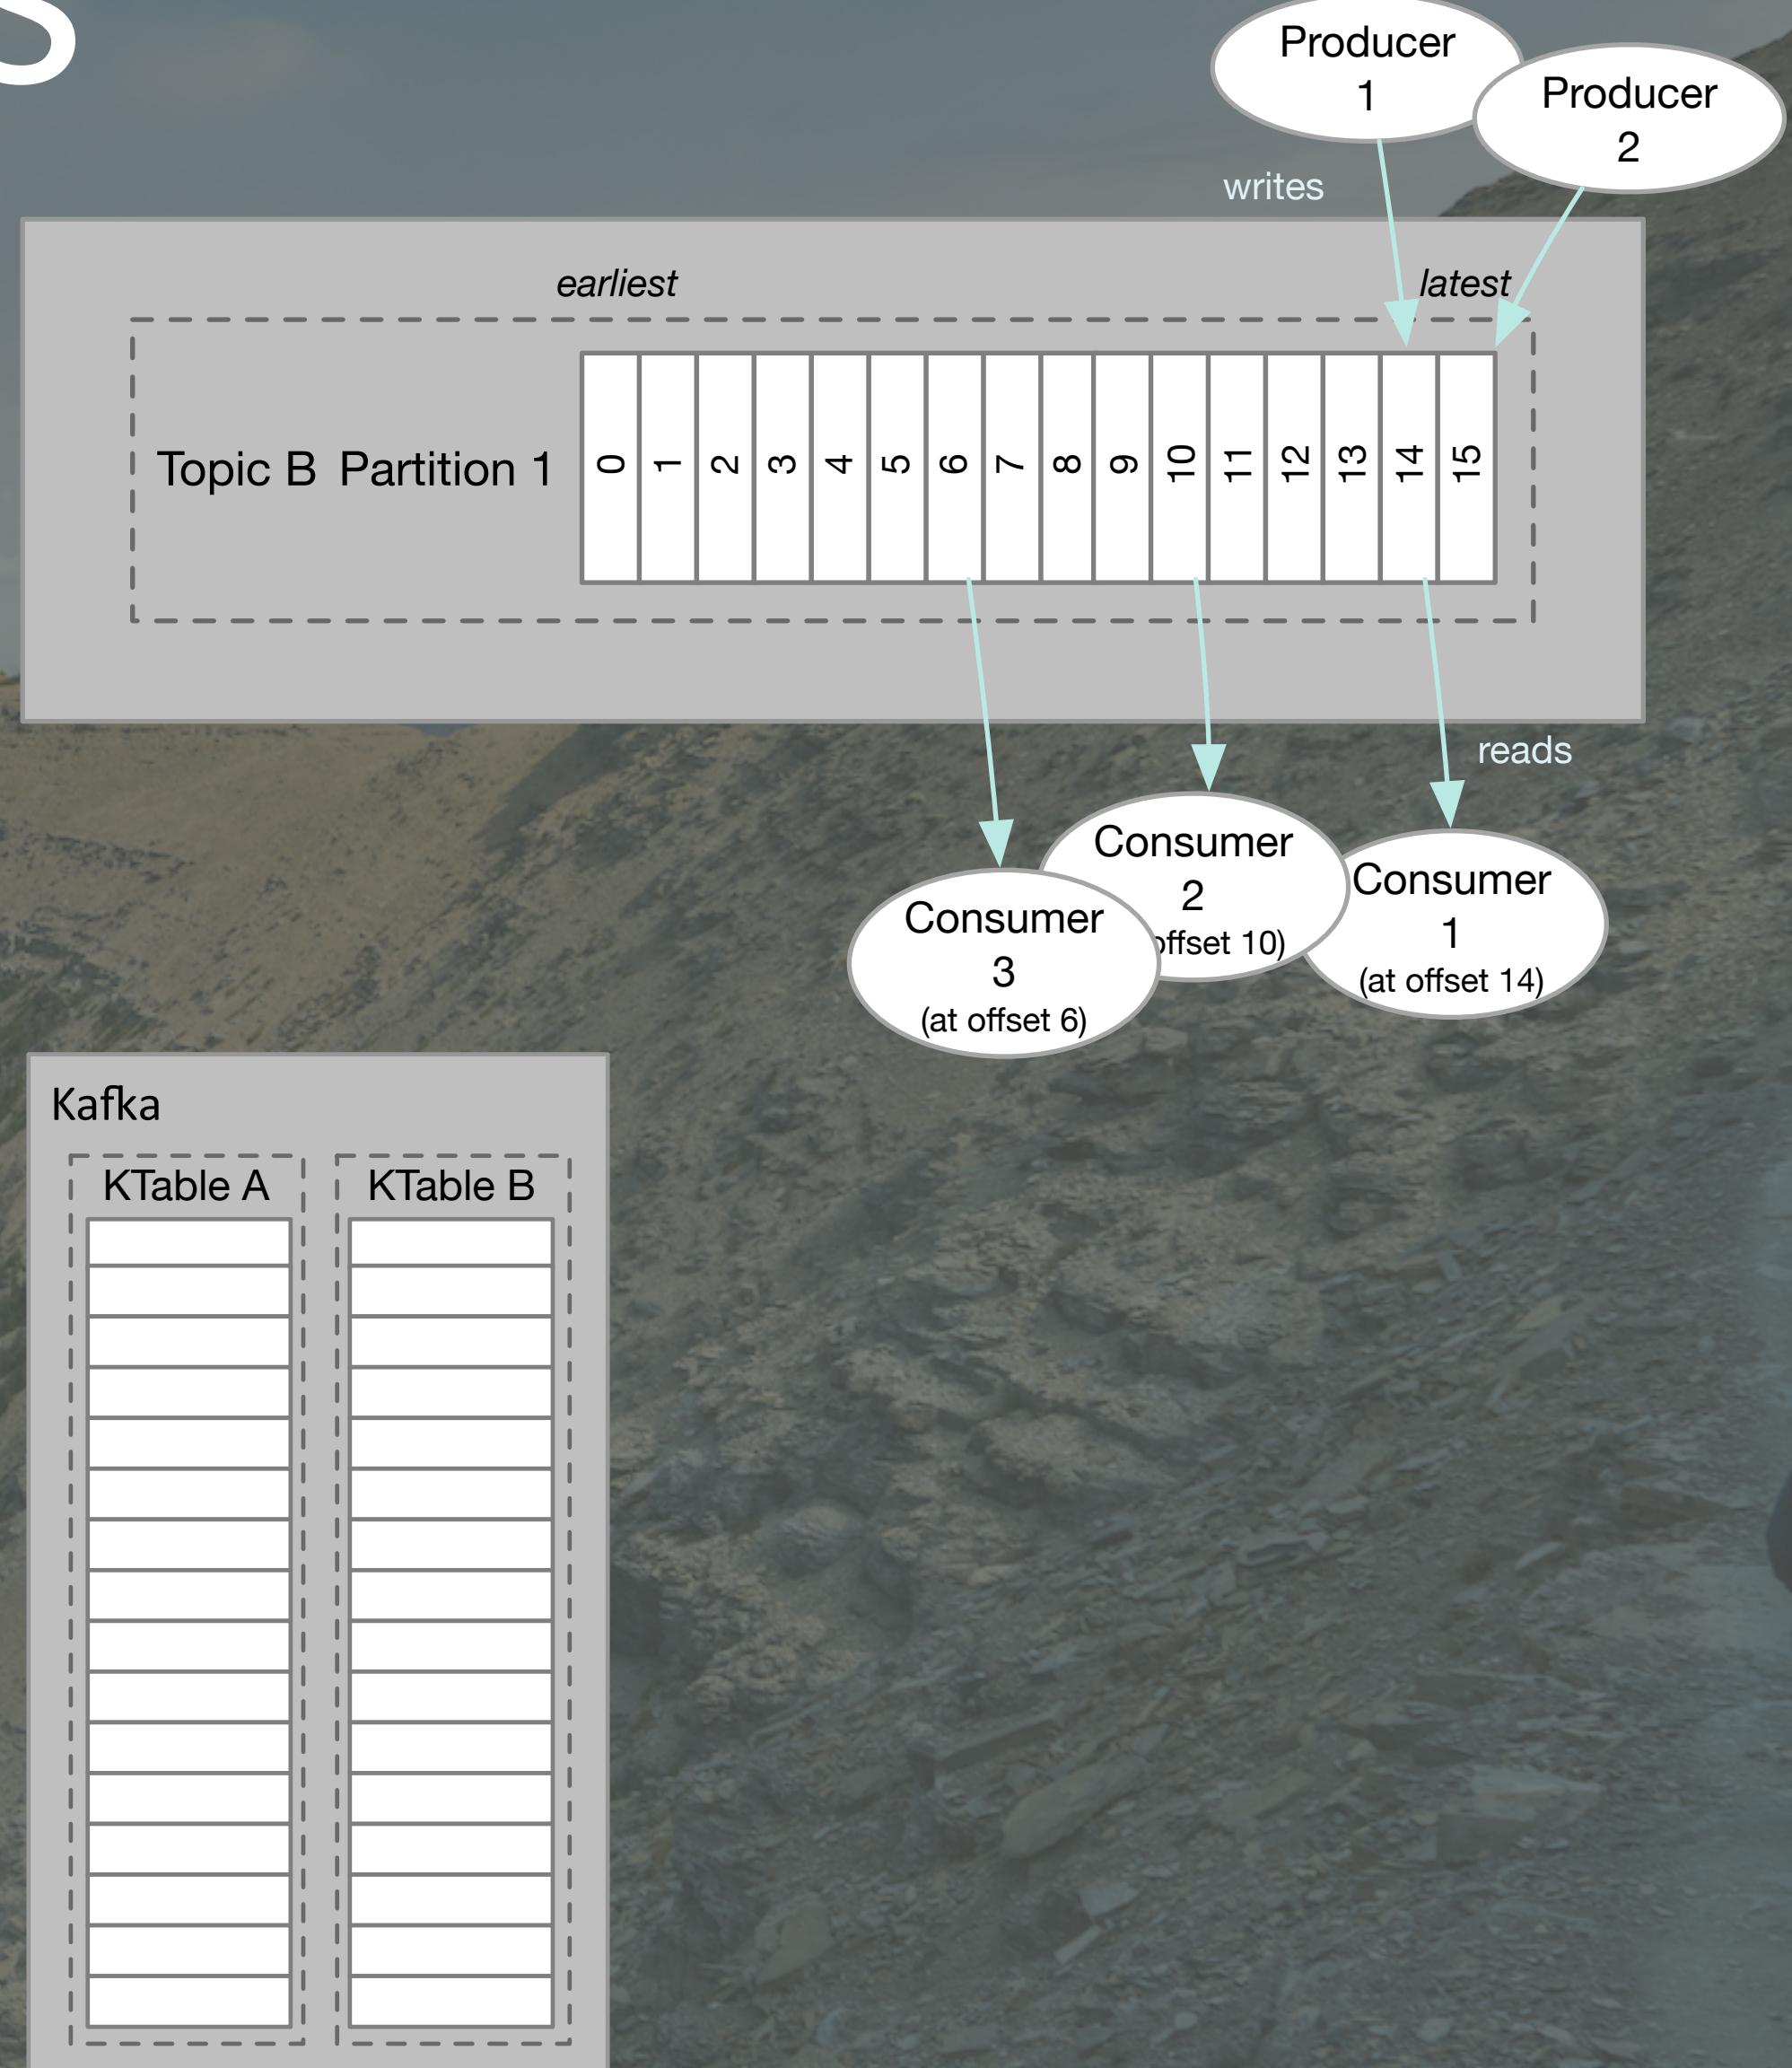
- Problem: I want to process each datum, one at a time.
- Problem: I want to track state (e.g., aggregations)



kafka

# Kafka Streams

- Solution: KStream
  - See the whole stream
- Solution: KTable
  - Last value per key





kafka

# Kafka Streams

- Problem: I need to scale horizontally
- Solution: Manually partition your Kafka topics



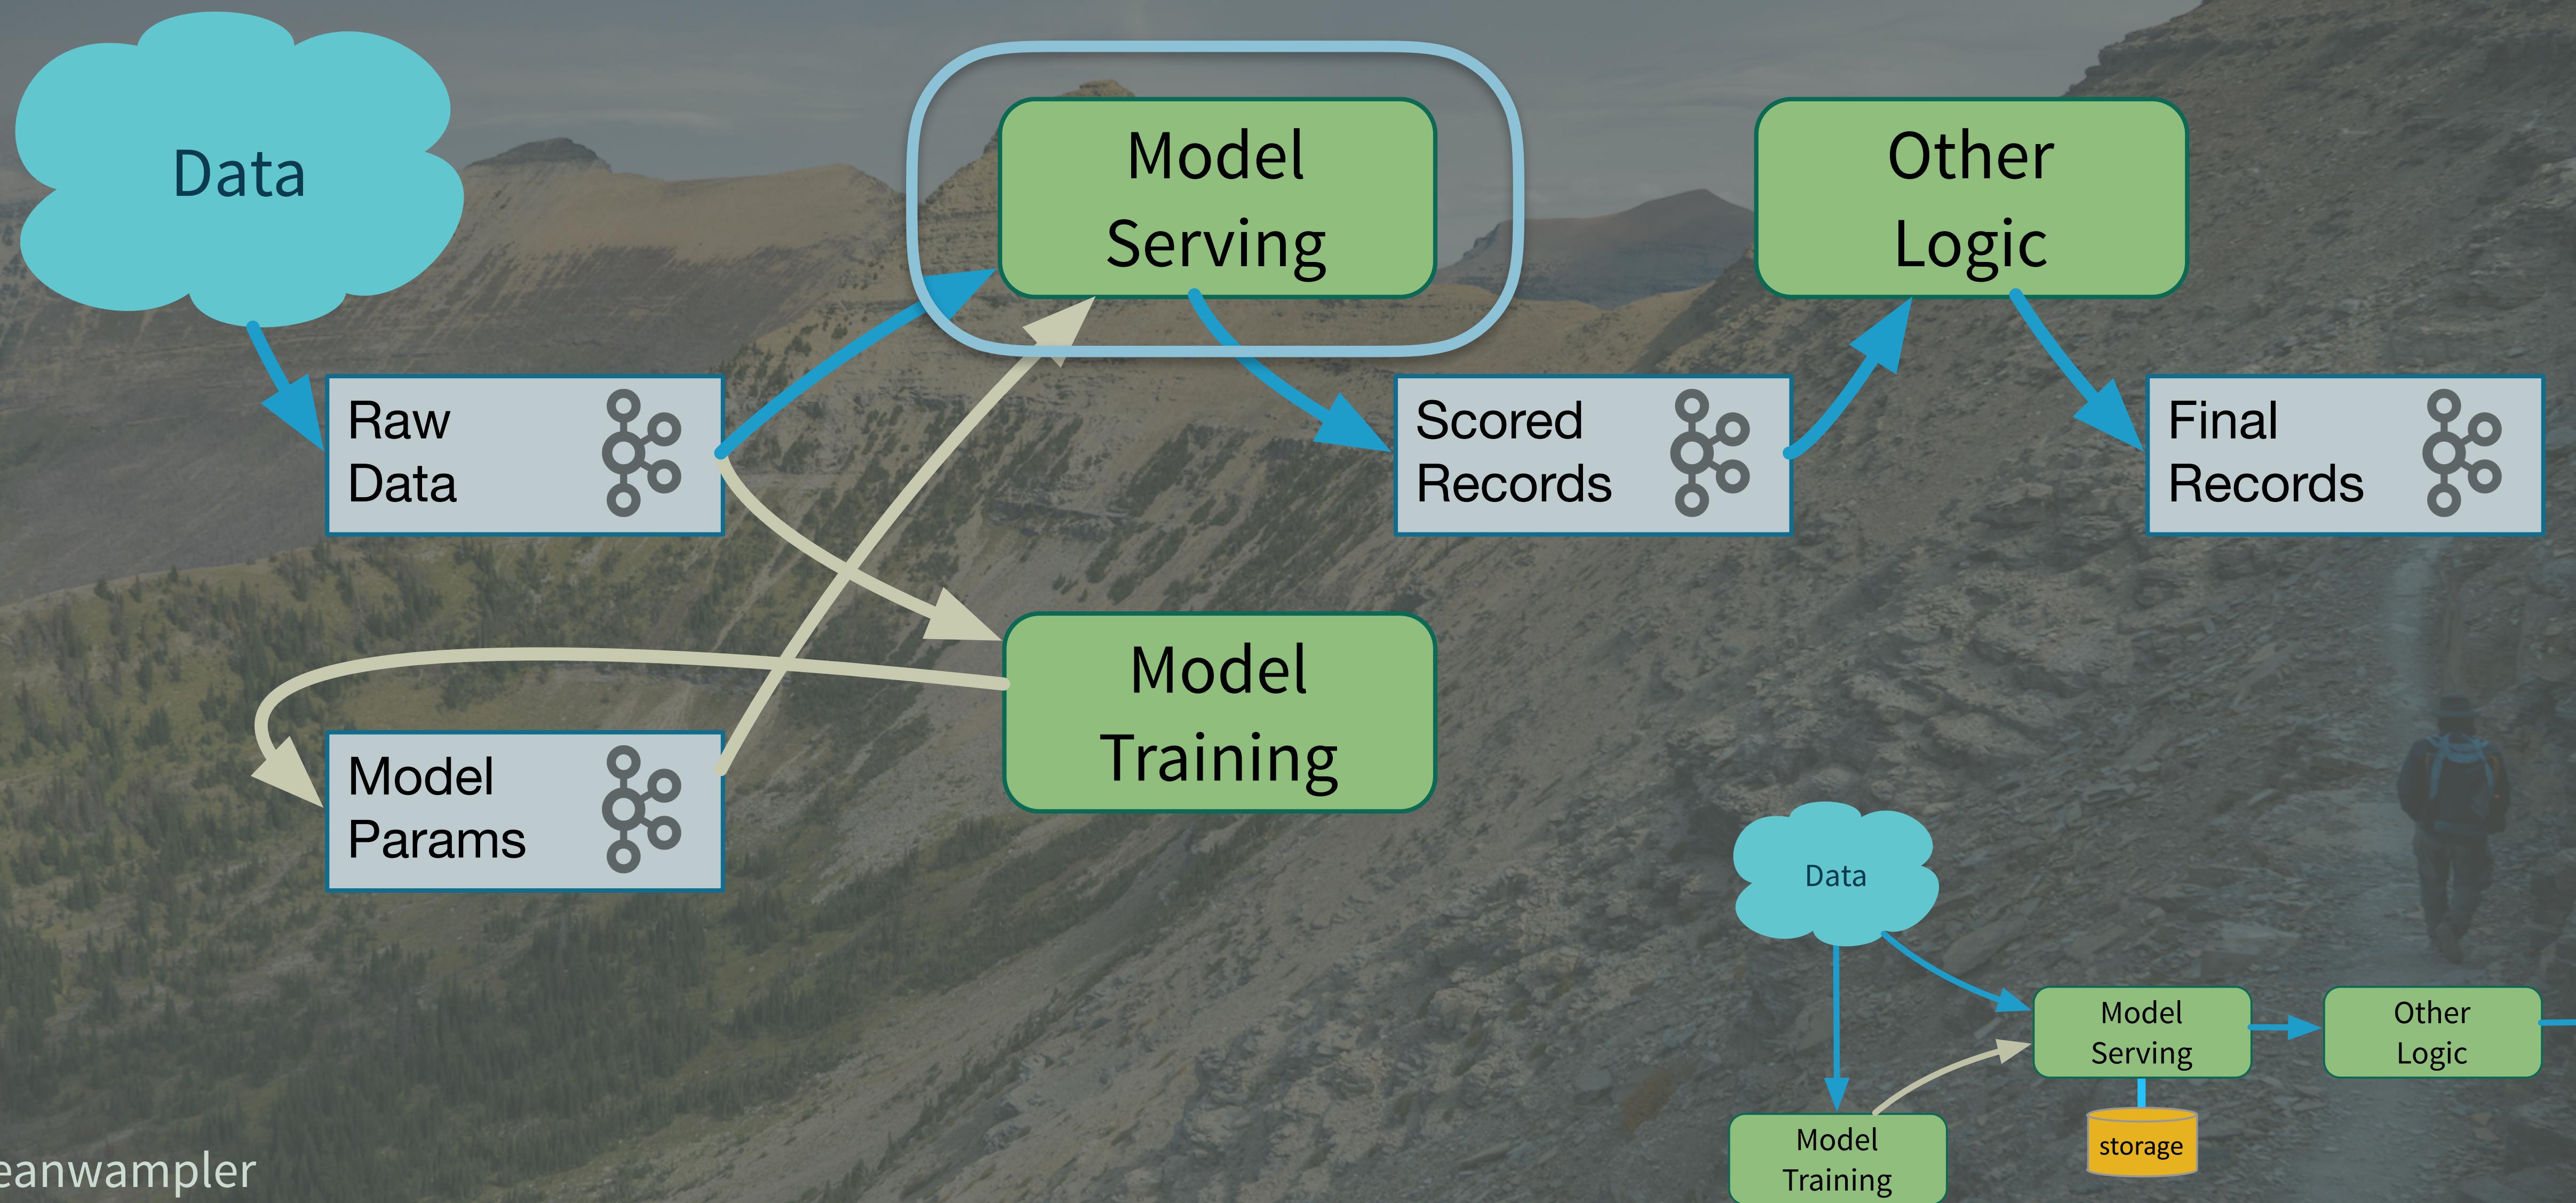
kafka

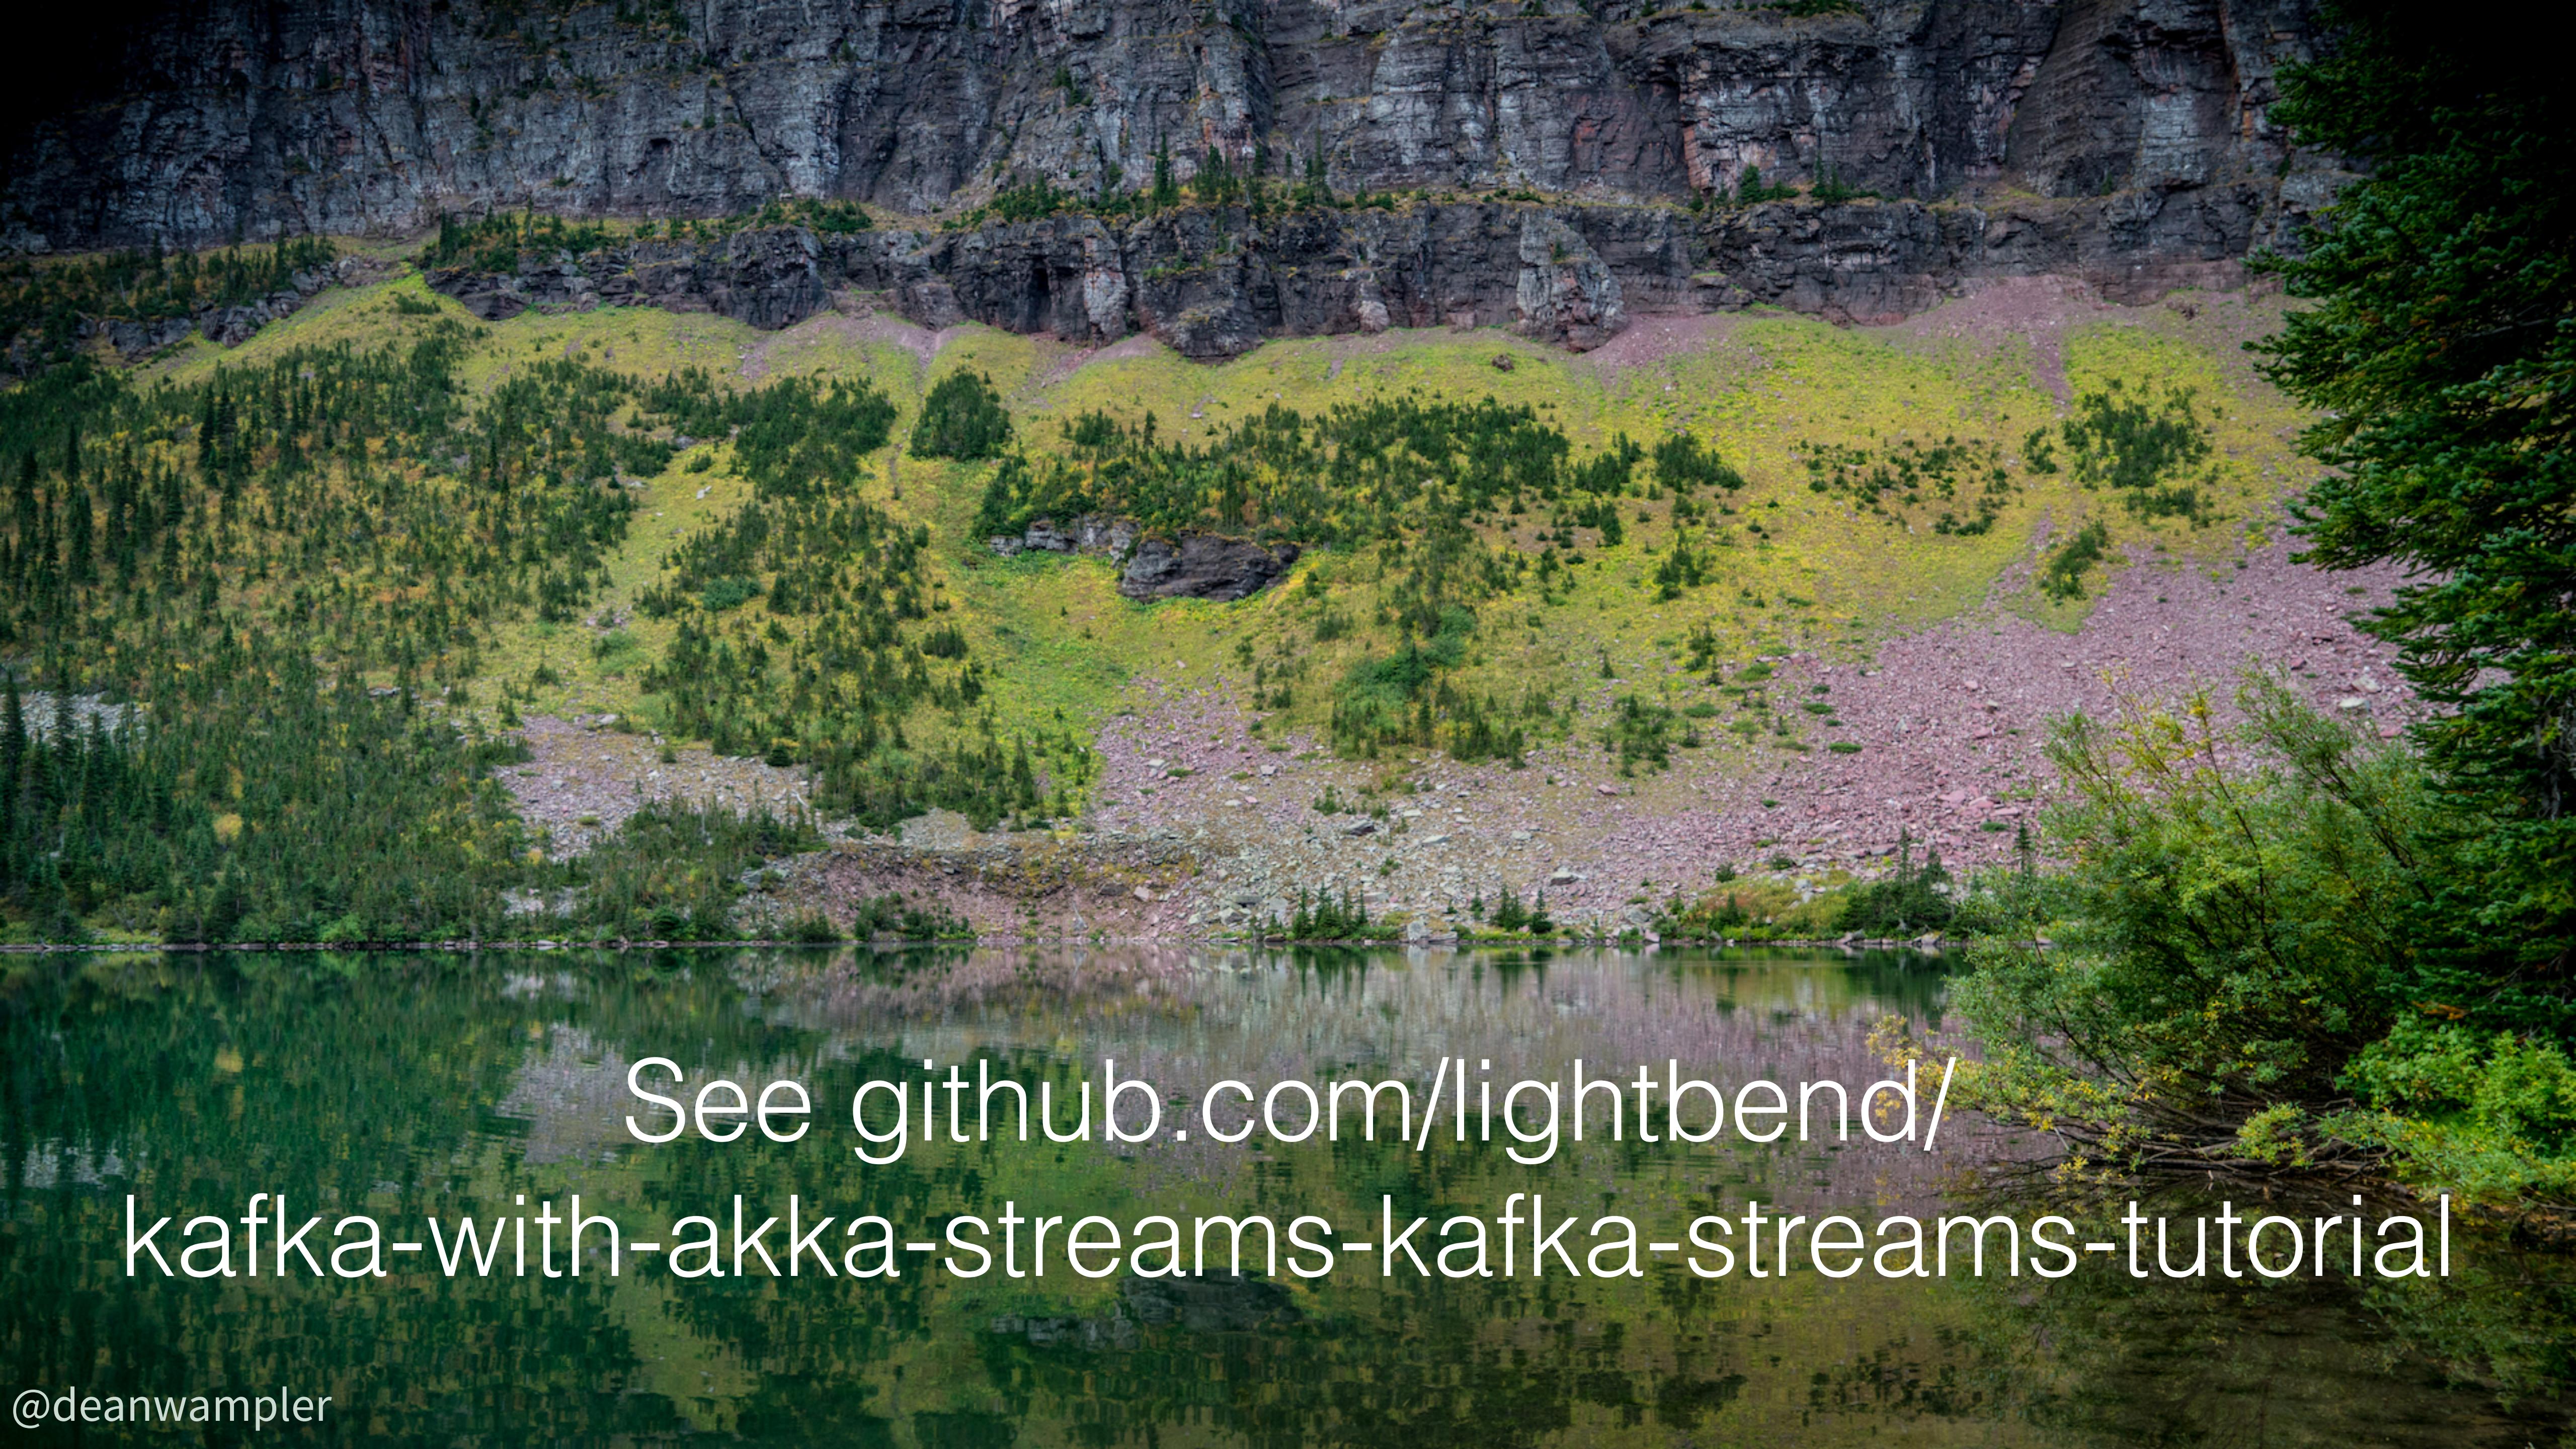
# Kafka Streams

- Java API
- Scala API: written by Lightbend
  - Debasish Ghosh, Boris Lublinsky, Sean Glover, et al.
  - If you know about *queryable state*, see
  - <https://github.com/lightbend/kafka-streams-query>
- SQL!!



# Kafka Streams Example



A scenic view of a lake surrounded by steep, rocky mountains with patches of green vegetation.

See [github.com/lightbend/  
kafka-with-akka-streams-kafka-streams-tutorial](https://github.com/lightbend/kafka-with-akka-streams-kafka-streams-tutorial)

```

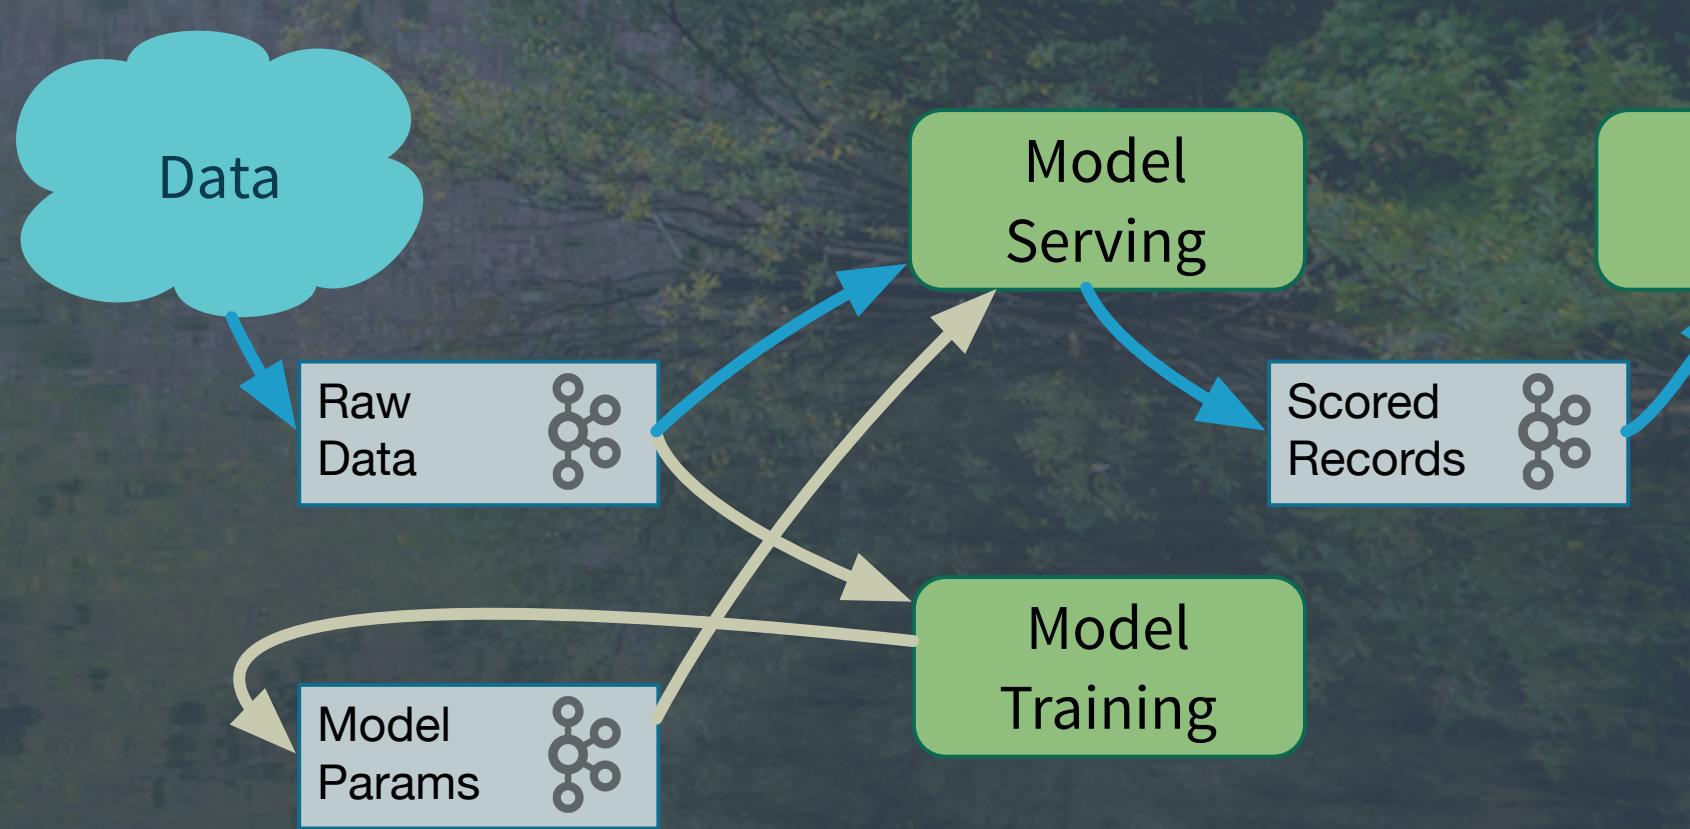
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```

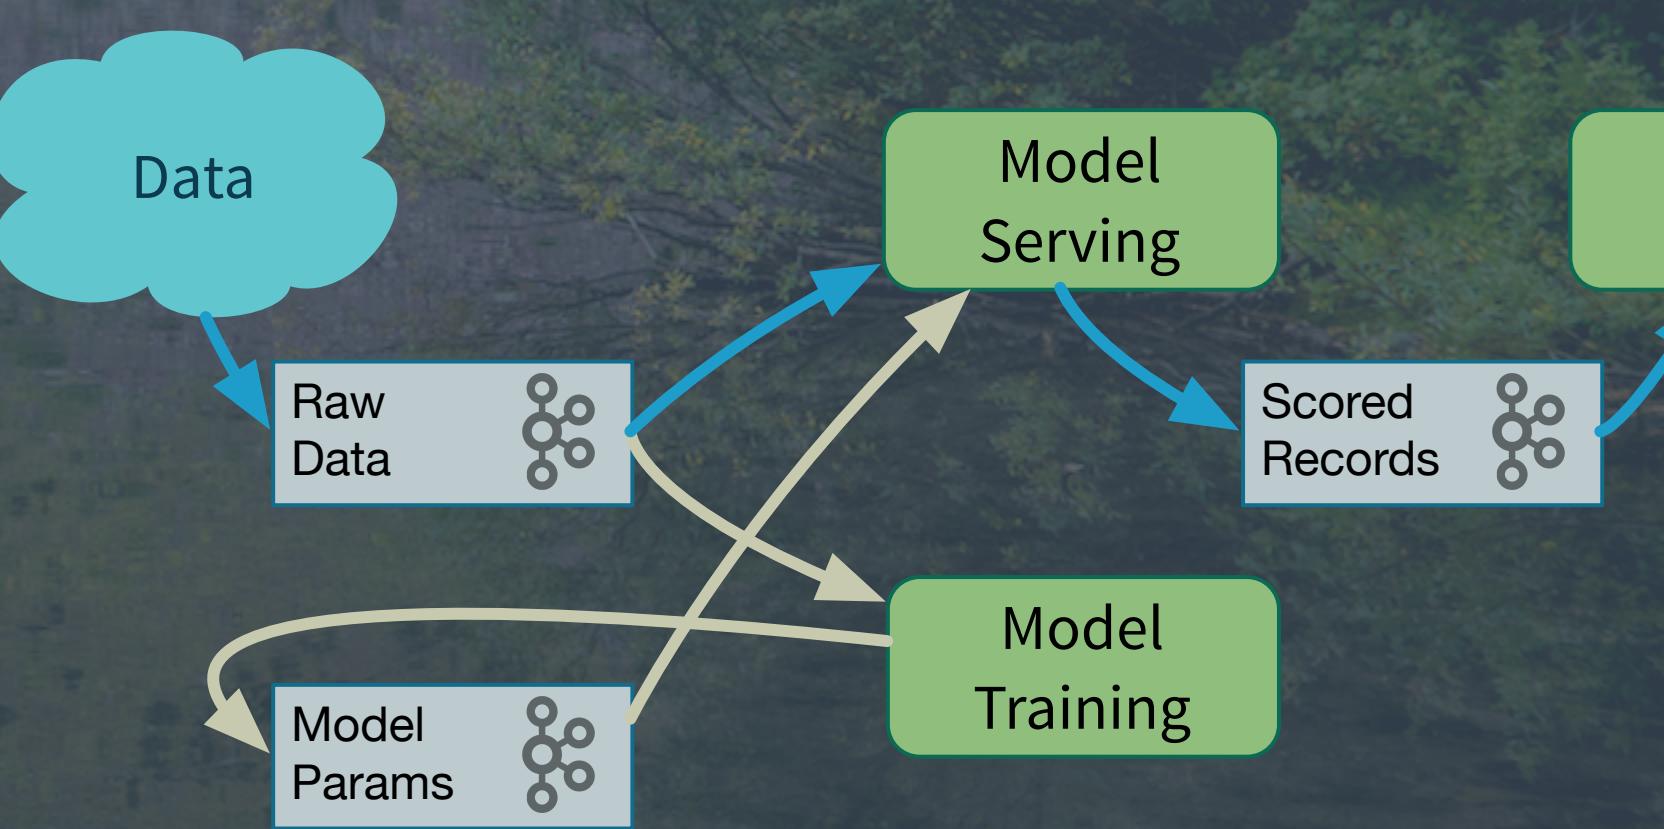


```
val builder = new StreamsBuilders // New Scala Wrapper API.
```

```
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)
```

```
val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())
```

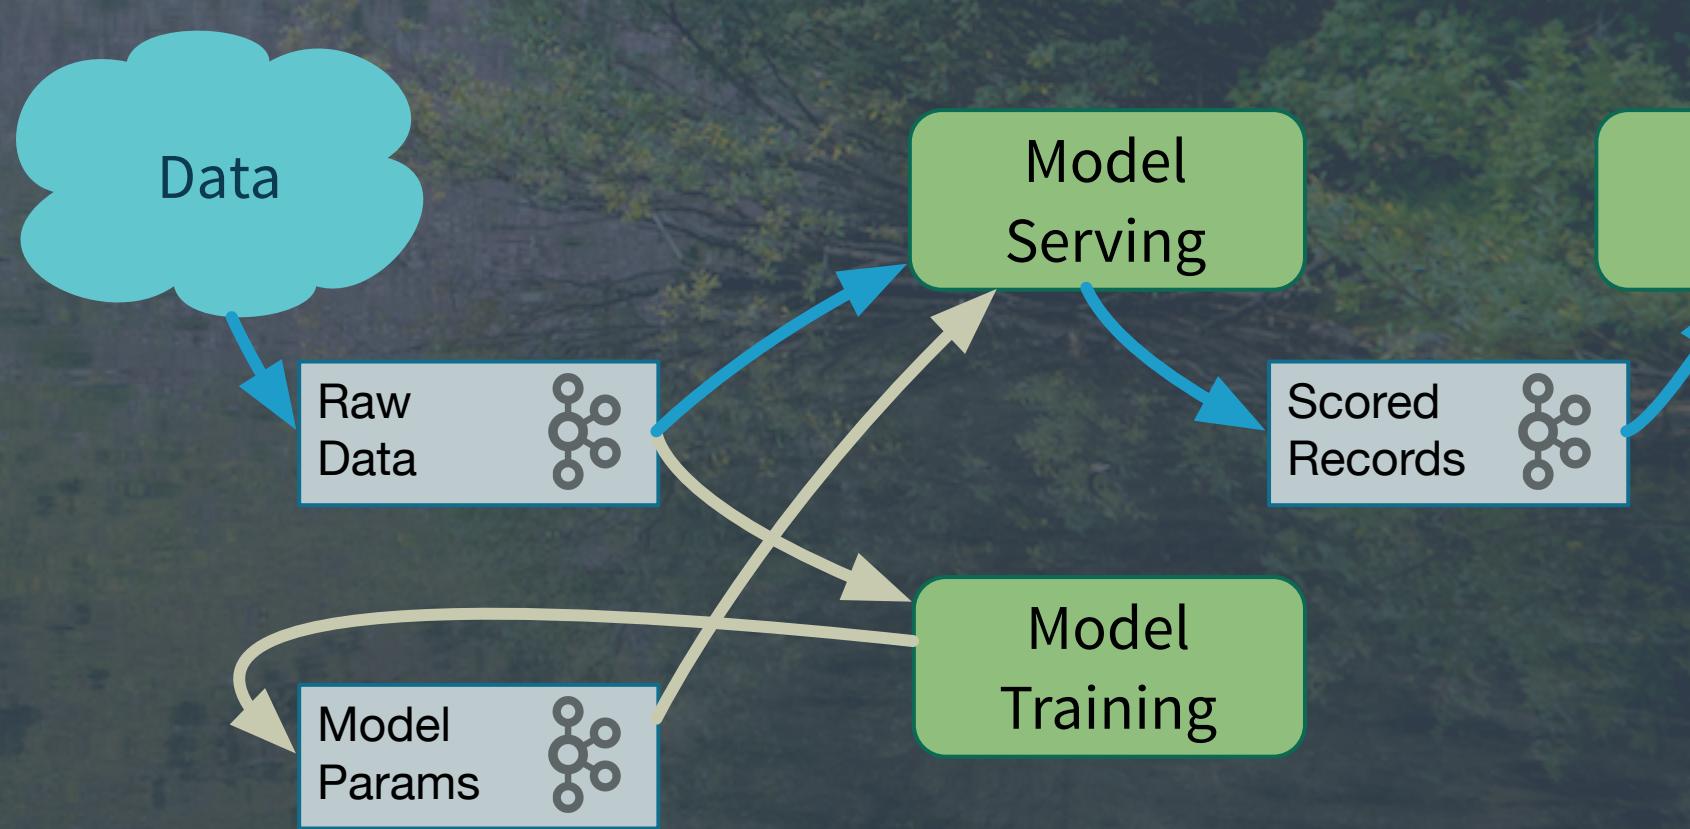


```
val builder = new StreamsBuilders // New Scala Wrapper API.
```

```
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(score(score, record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())
```



```

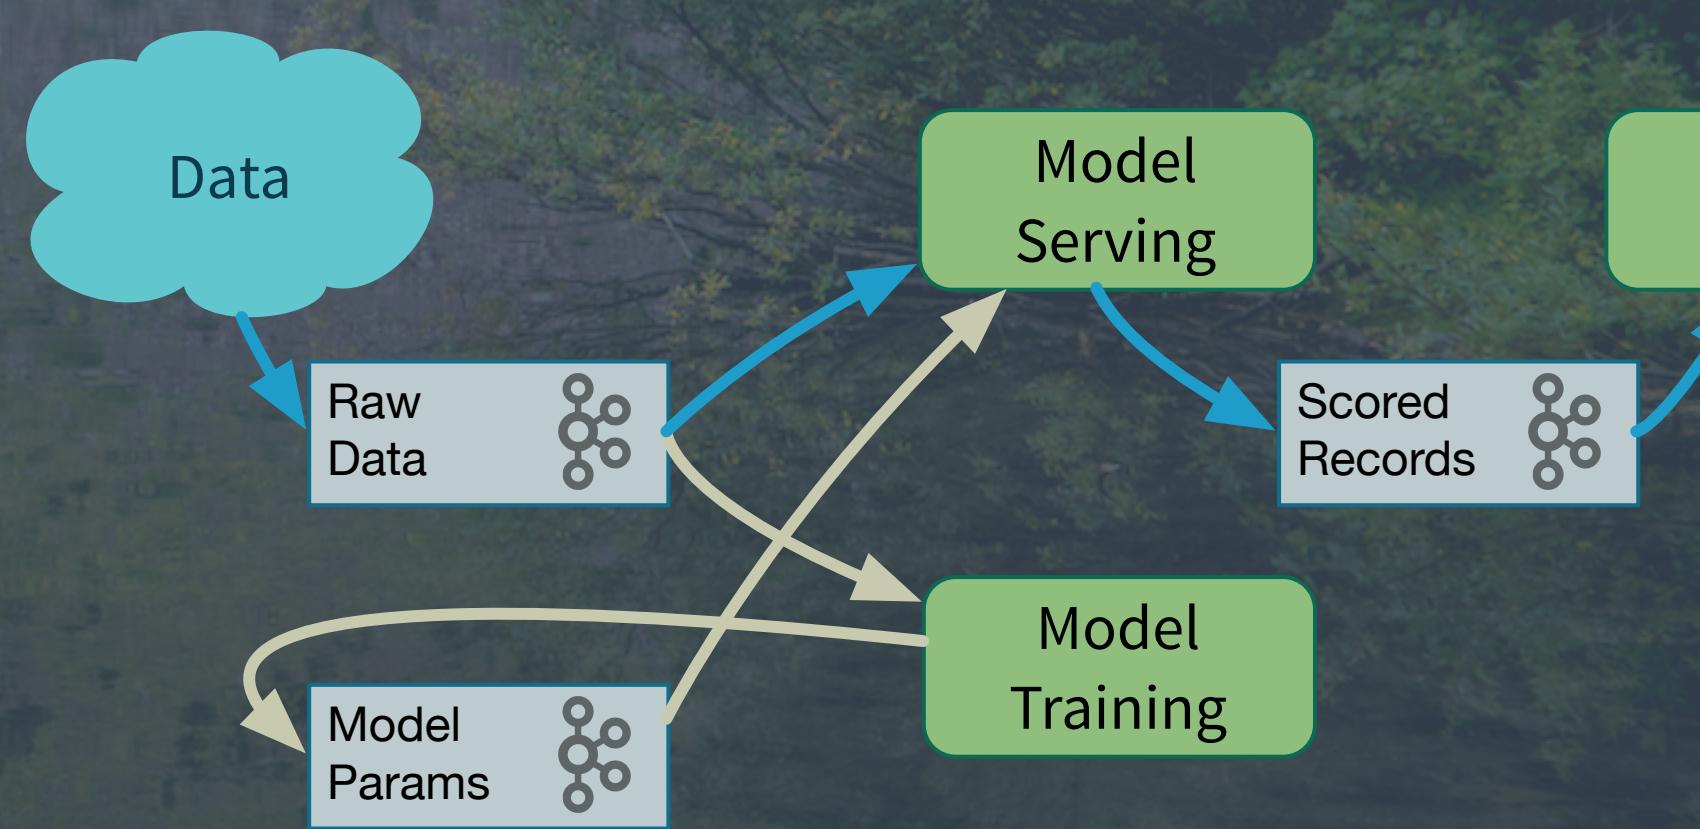
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(scorer.score(record), record))
  .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```



```

val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

```

```

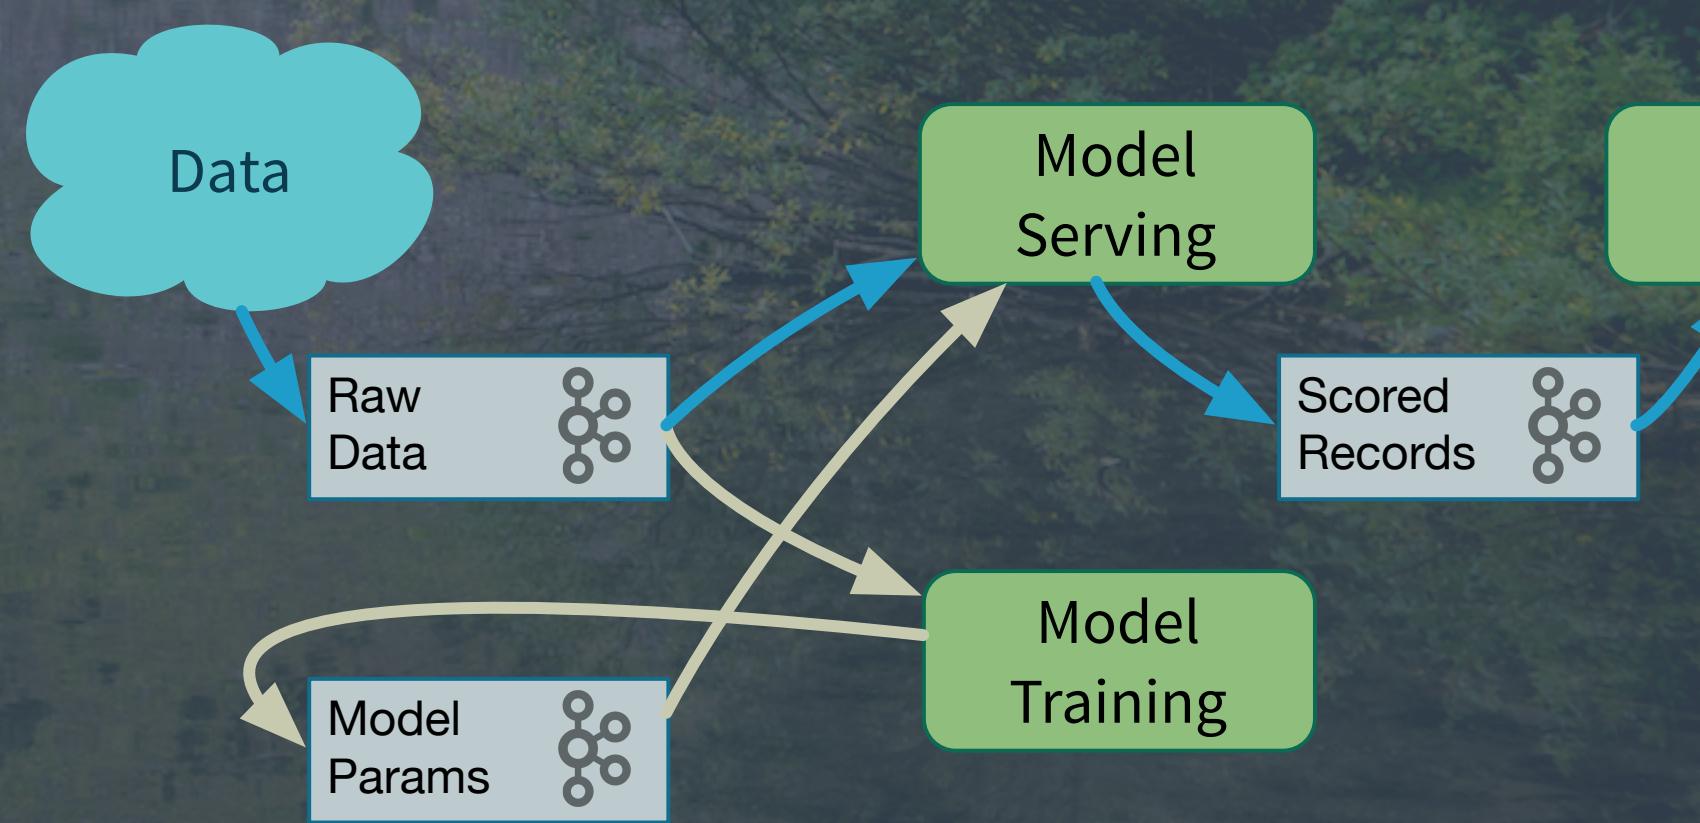
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() => modelProcessor, ...) // Set up actual model
data.mapValues(bytes => DataRecord.parseBytes(bytes))
  .filter((key, record) => record.valid)
  .mapValues(record => new ScoredRecord(score(score(record), record)))
  .to(scoredRecordsTopic)

```

```

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

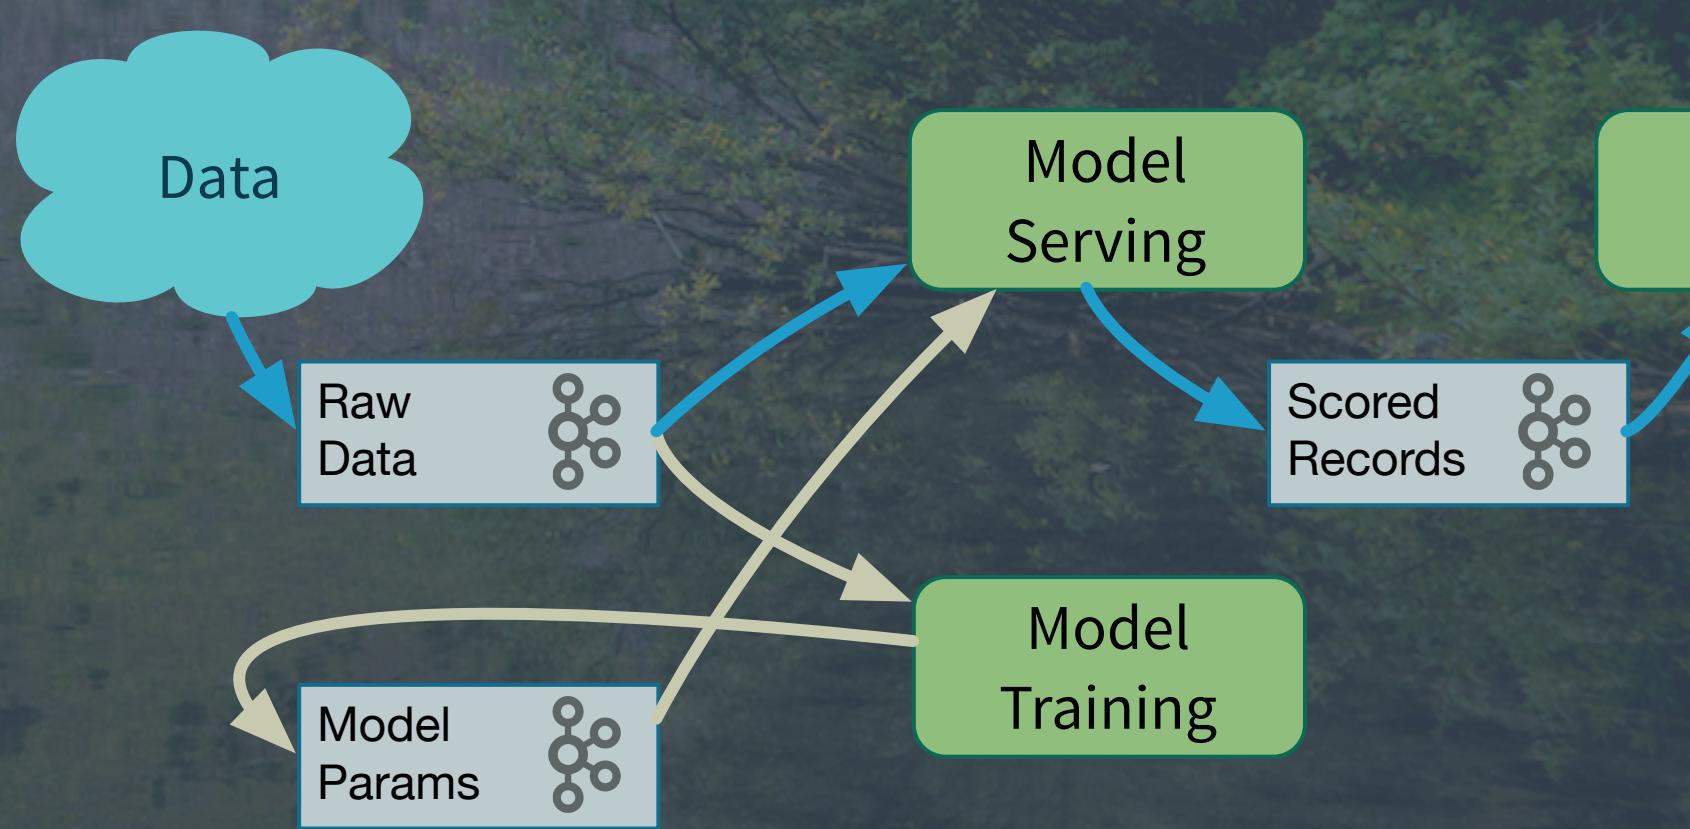
```



```
val builder = new StreamsBuilders // New Scala Wrapper API.  
  
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)  
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)  
val modelProcessor = new ModelProcessor  
val scorer = new Scorer(modelProcessor) // scorer.score(record) used
```

```
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record  
.filter((key, model) => model.valid) // Successful?  
.mapValues(model => ModelImpl.findModel(model))  
.process(() => modelProcessor, ...) // Set up actual model  
data.mapValues(bytes => DataRecord.parseBytes(bytes))  
.filter((key, record) => record.valid)  
.mapValues(record => new ScoredRecord(score(score(record), record)))  
.to(scoredRecordsTopic)
```

```
val streams = new KafkaStreams(  
  builder.build, streamsConfiguration)  
streams.start()  
sys.addShutdownHook(streams.close())
```



```

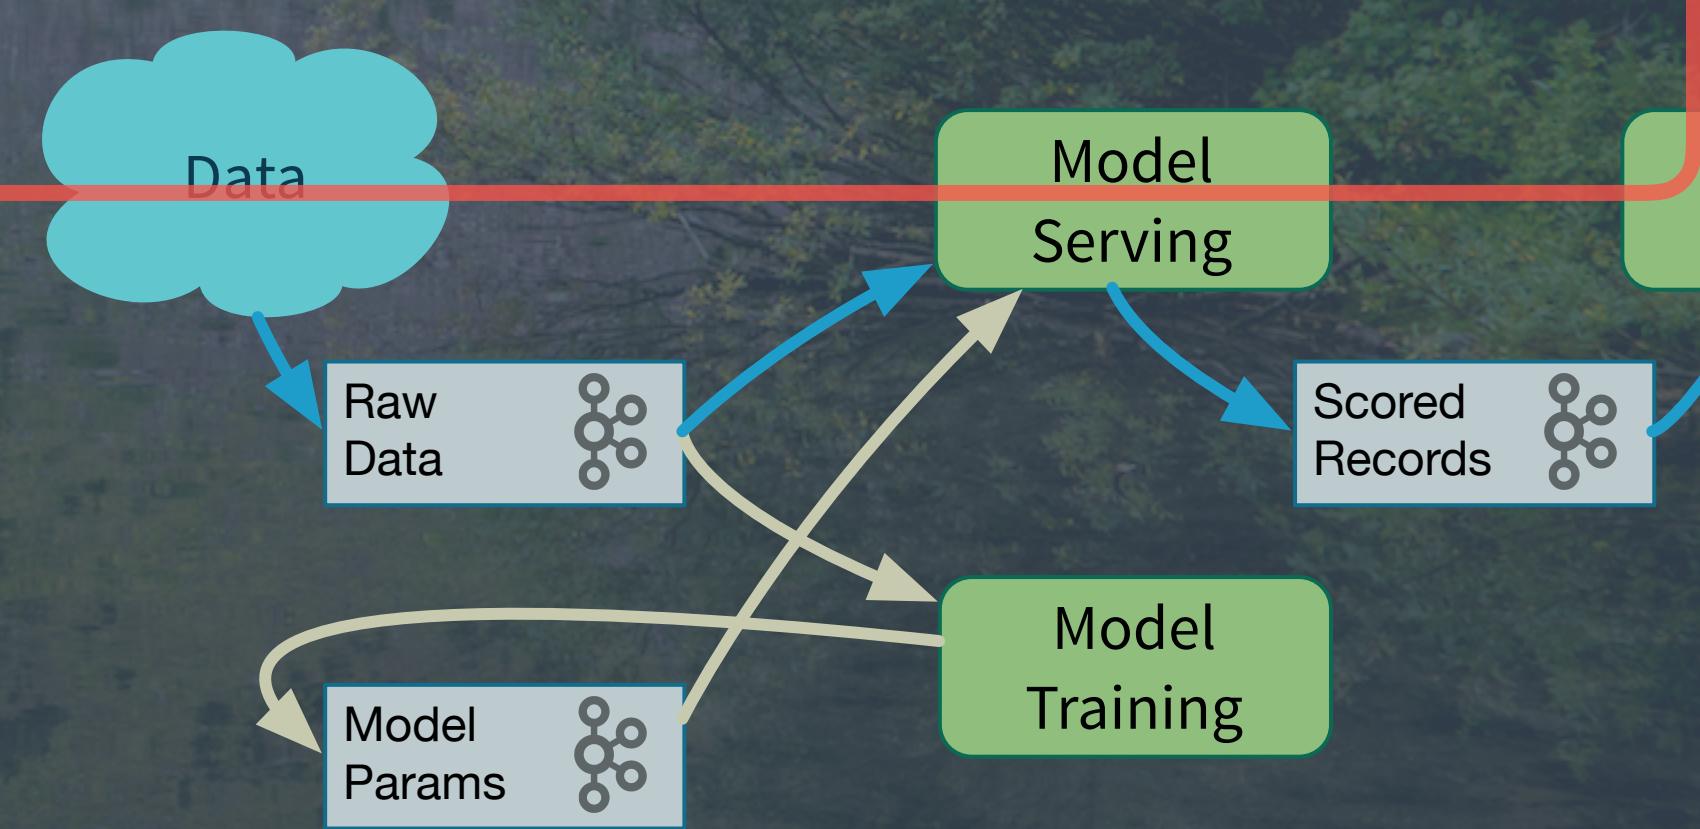
val builder = new StreamsBuilders // New Scala Wrapper API.

val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)
val modelProcessor = new ModelProcessor
val scorer = new Scorer(modelProcessor) // scorer.score(record) used

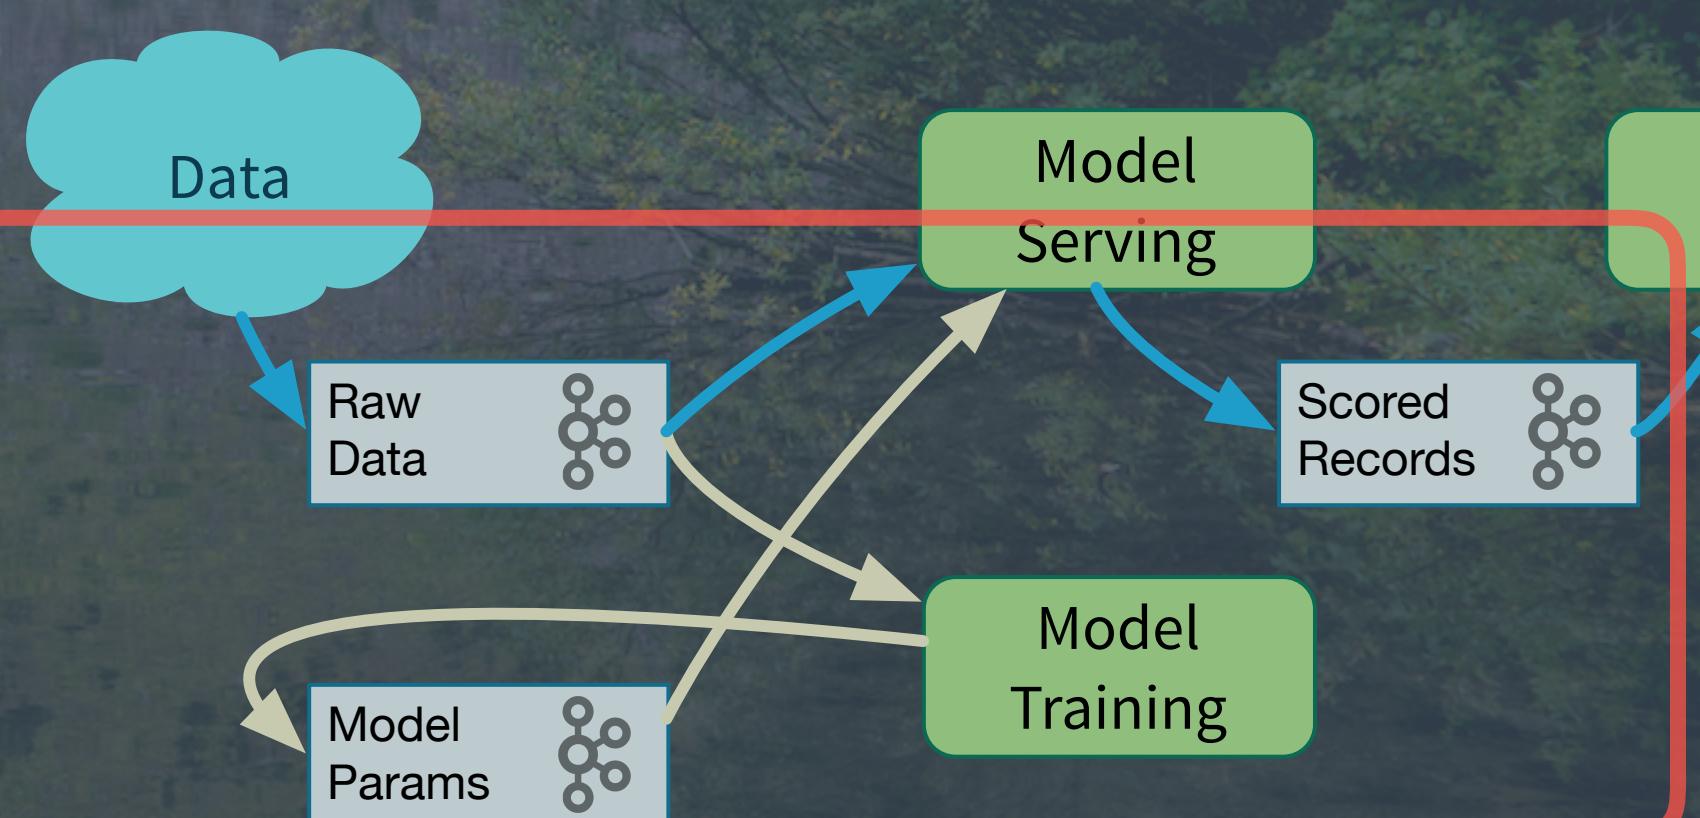
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record
  .filter((key, model) => model.valid) // Successful?
  .mapValues(model => ModelImpl.findModel(model))
  .process(() -> modelProcessor, ...) // Set up actual model
  data.mapValues(bytes => DataRecord.parseBytes(bytes))
    .filter((key, record) => record.valid)
    .mapValues(record => new ScoredRecord(score(score(record), record)))
    .to(scoredRecordsTopic)

val streams = new KafkaStreams(
  builder.build, streamsConfiguration)
streams.start()
sys.addShutdownHook(streams.close())

```



```
val builder = new StreamsBuilders // New Scala Wrapper API.  
  
val data  = builder.stream[Array[Byte], Array[Byte]](rawDataTopic)  
val model = builder.stream[Array[Byte], Array[Byte]](modelTopic)  
val modelProcessor = new ModelProcessor  
val scorer = new Scorer(modelProcessor) // scorer.score(record) used  
  
model.mapValues(bytes => Model.parseBytes(bytes)) // array => record  
.filter((key, model) => model.valid) // Successful?  
.mapValues(model => ModelImpl.findModel(model))  
.process(() => modelProcessor, ...) // Set up actual model  
data.mapValues(bytes => DataRecord.parseBytes(bytes))  
.filter((key, record) => record.valid)  
.mapValues(record => new ScoredRecord(score(score(record), record)))  
.to(scoredRecordsTopic)  
  
val streams = new KafkaStreams(  
  builder.build, streamsConfiguration)  
streams.start()  
sys.addShutdownHook(streams.close())
```



# What's Missing?

The rest of the microservice tools  
you need...

Embed your Kafka Streams code  
in microservices written with Akka  
or other JVM tools.

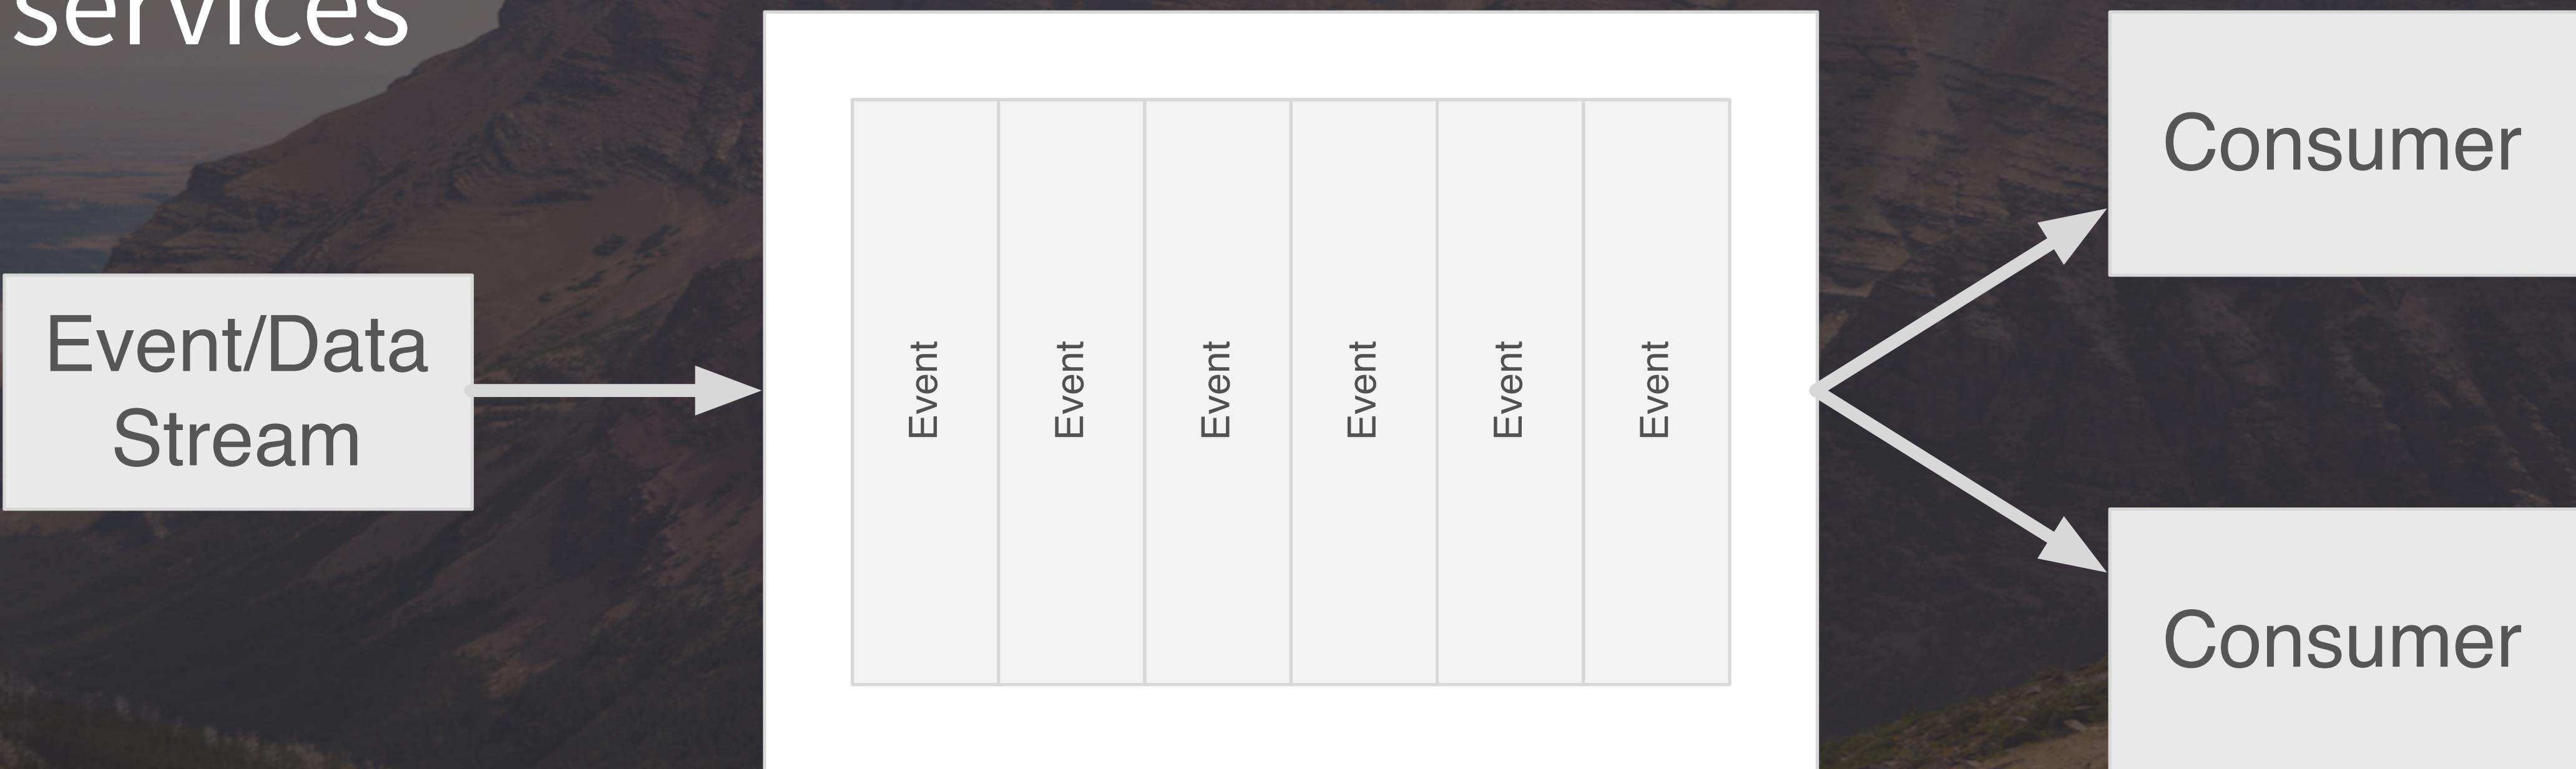


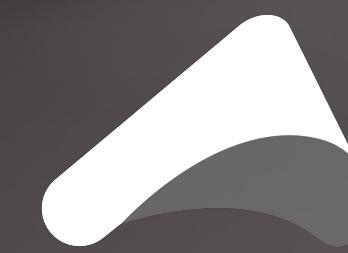
# akka Akka Streams



# Akka Streams

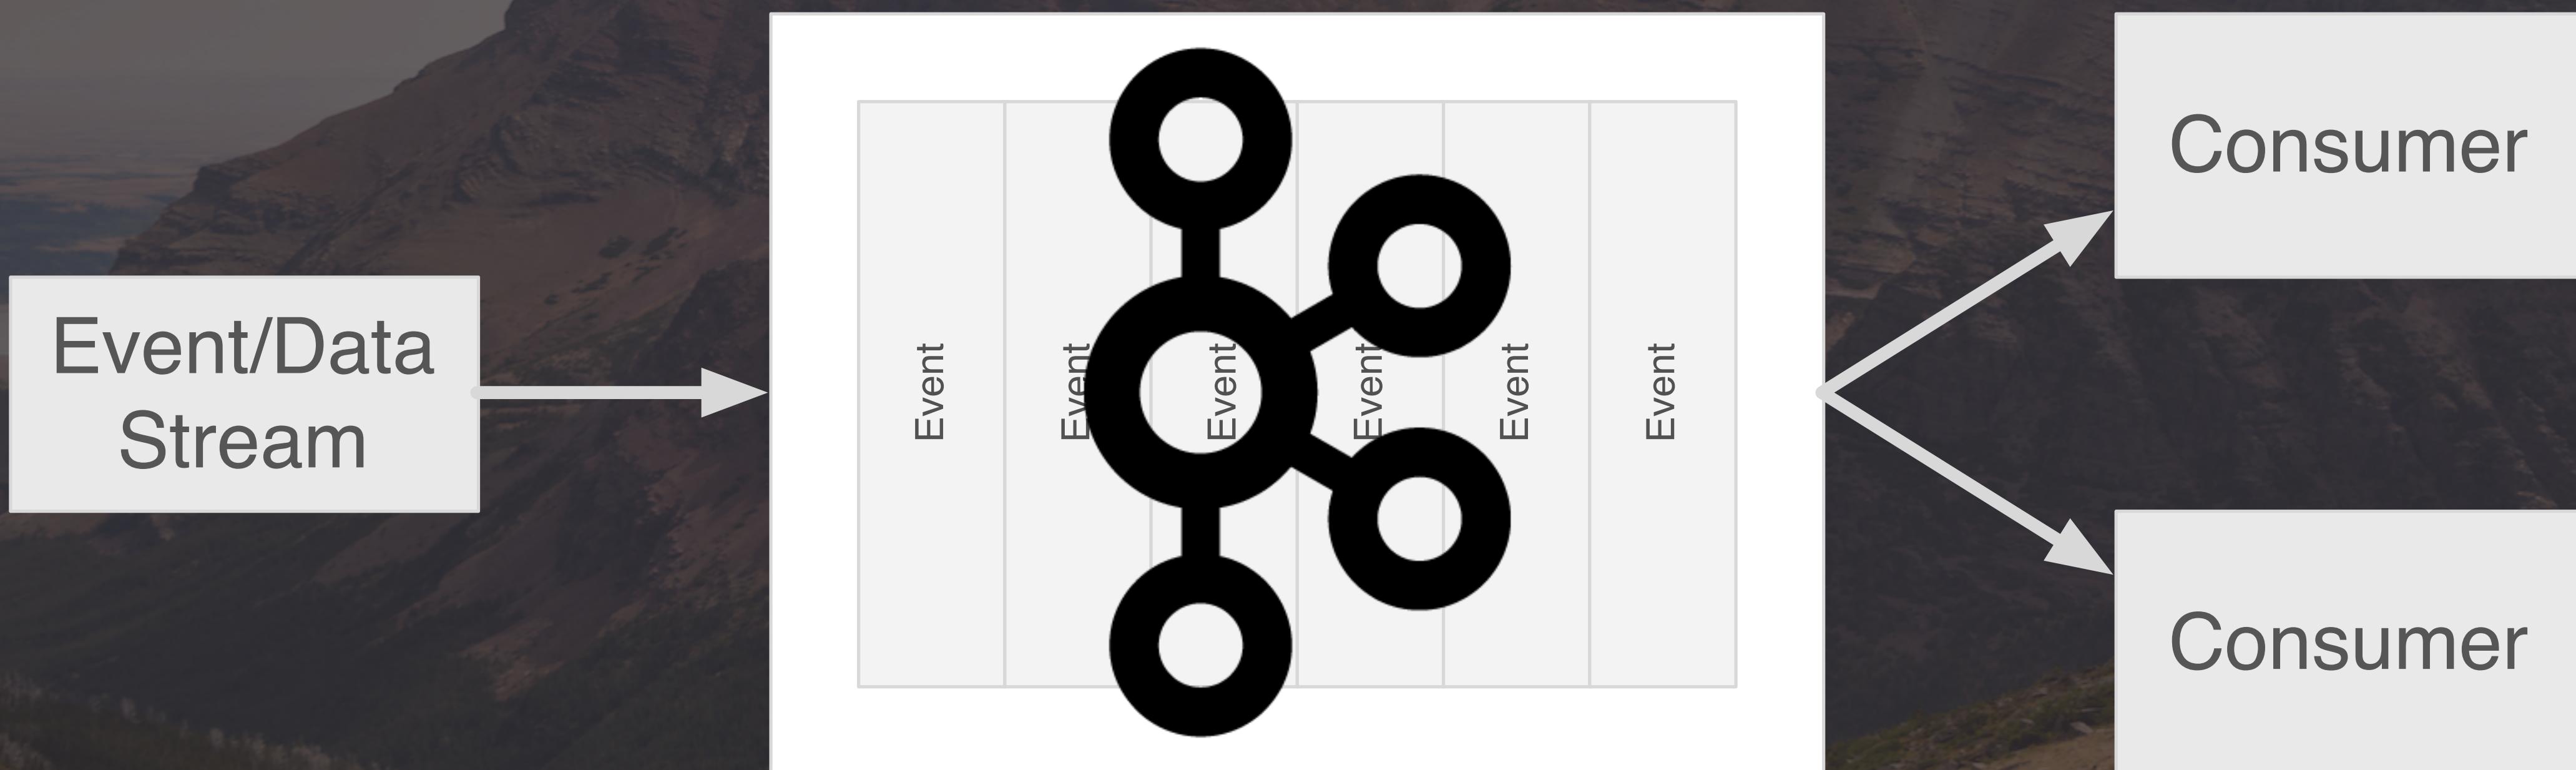
- Problem: I need reliable *buffering* between services

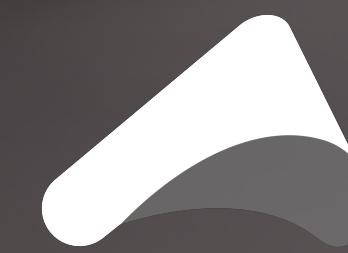




# Akka Streams

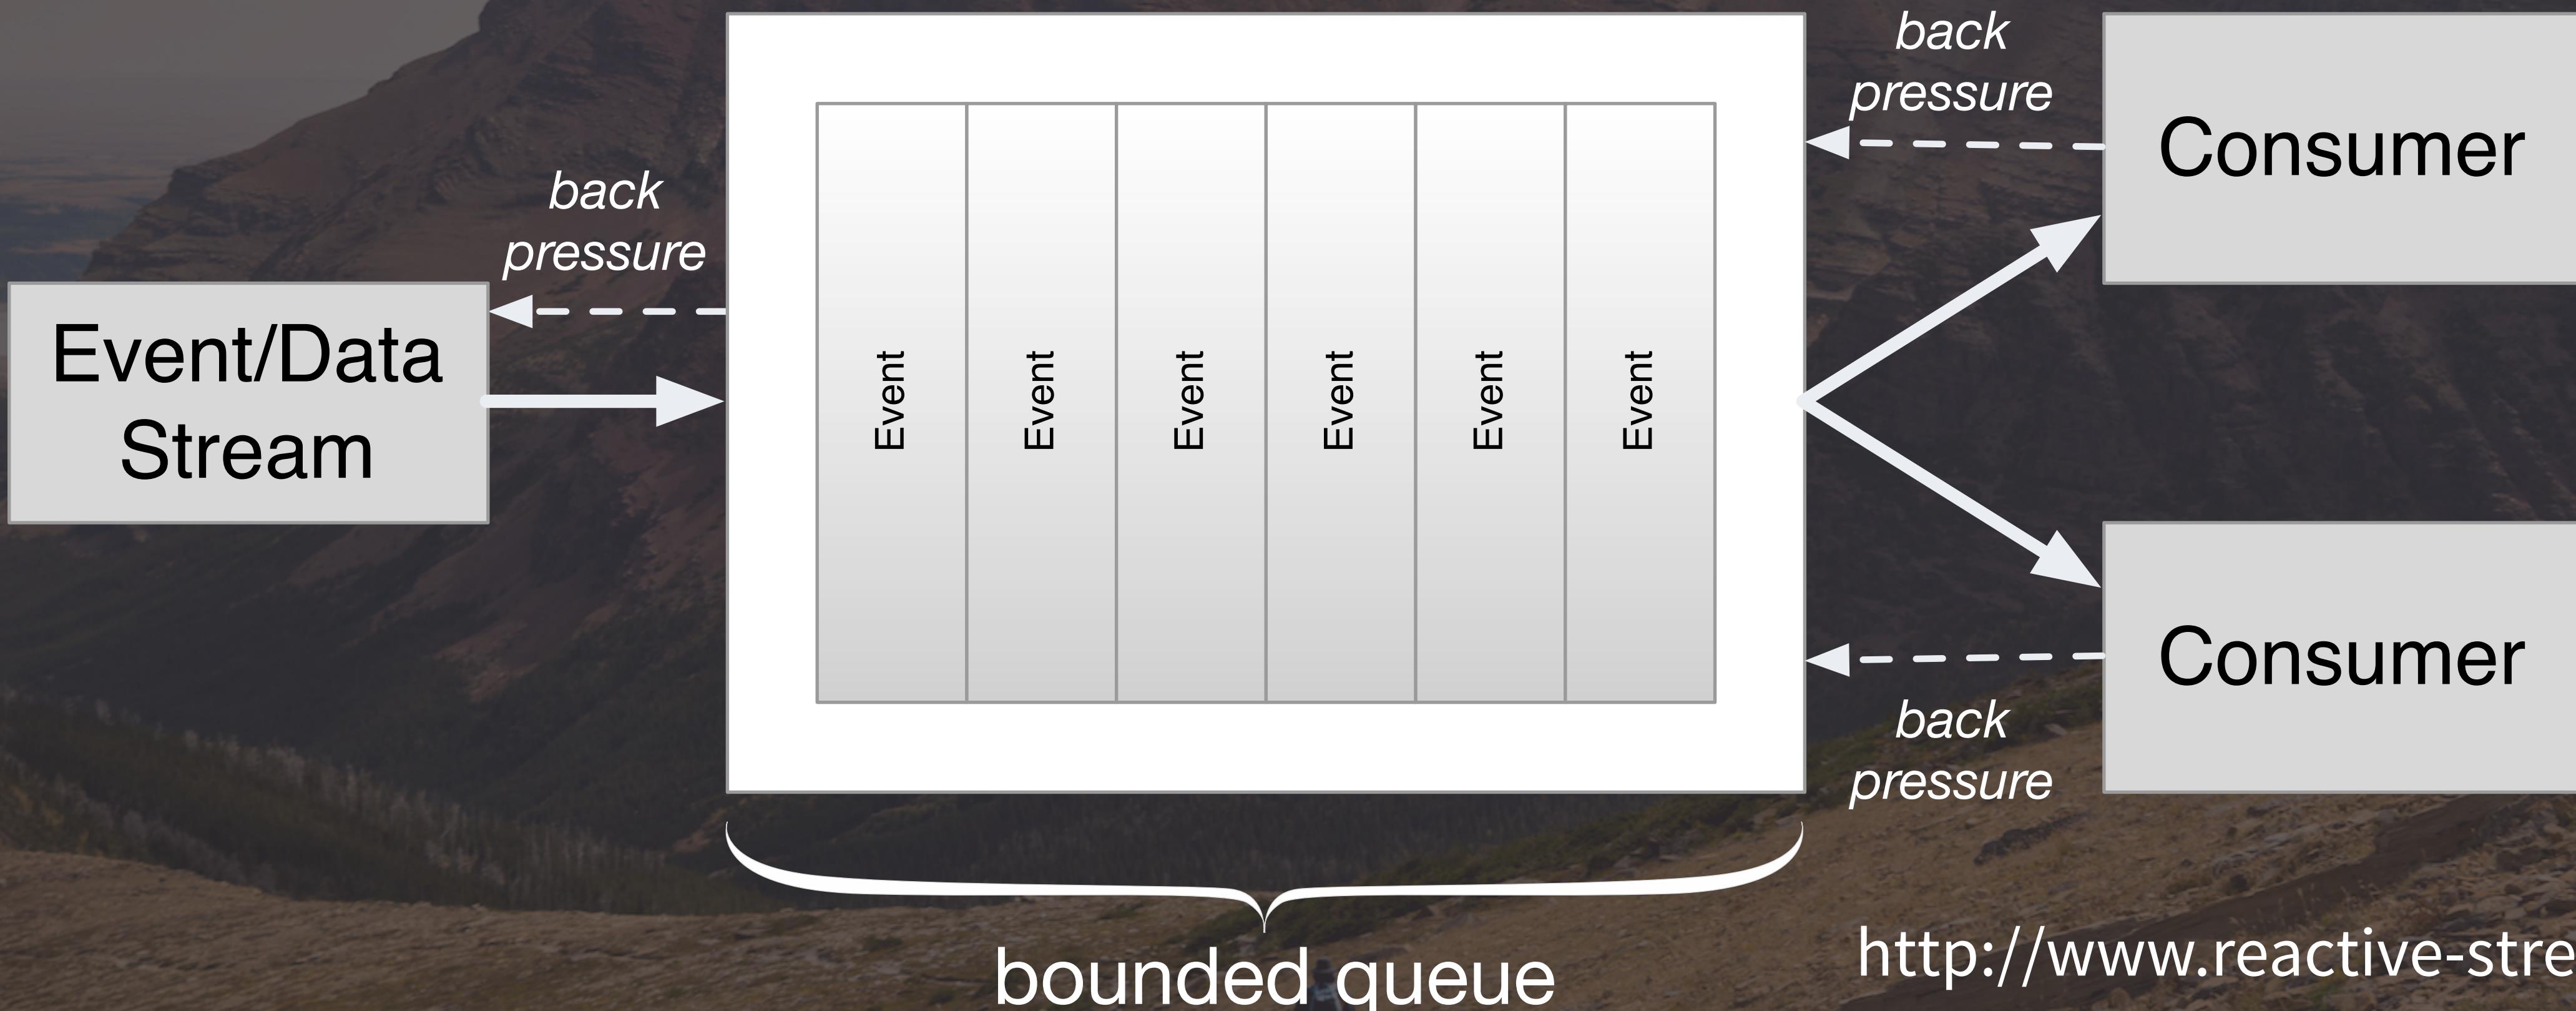
- Solution: Use a Kafka Topic!



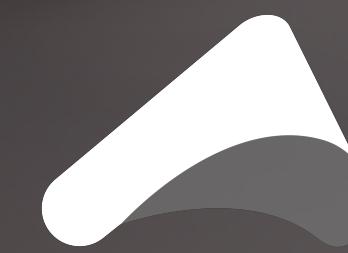


# Akka Streams

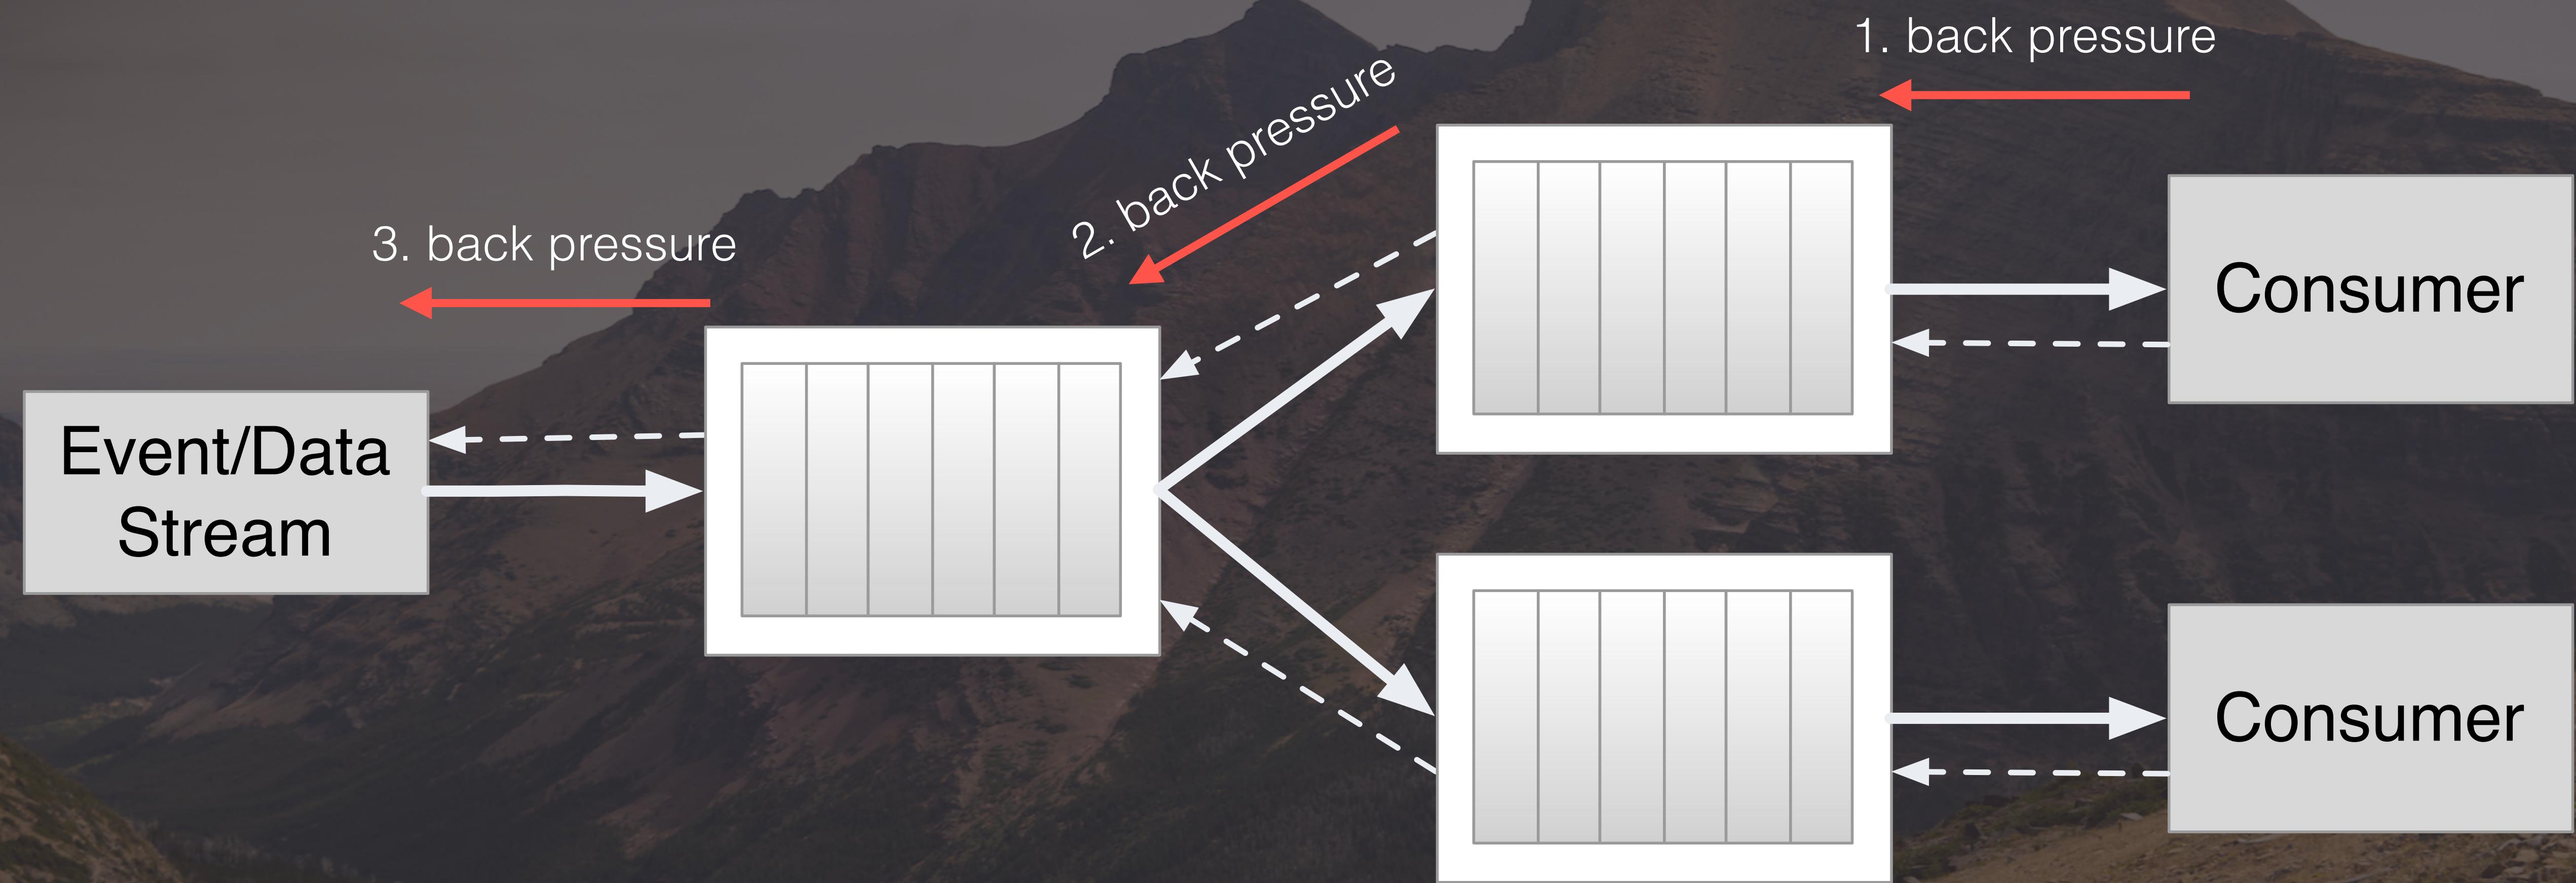
- Solution: Use *Reactive Streams* for flow control



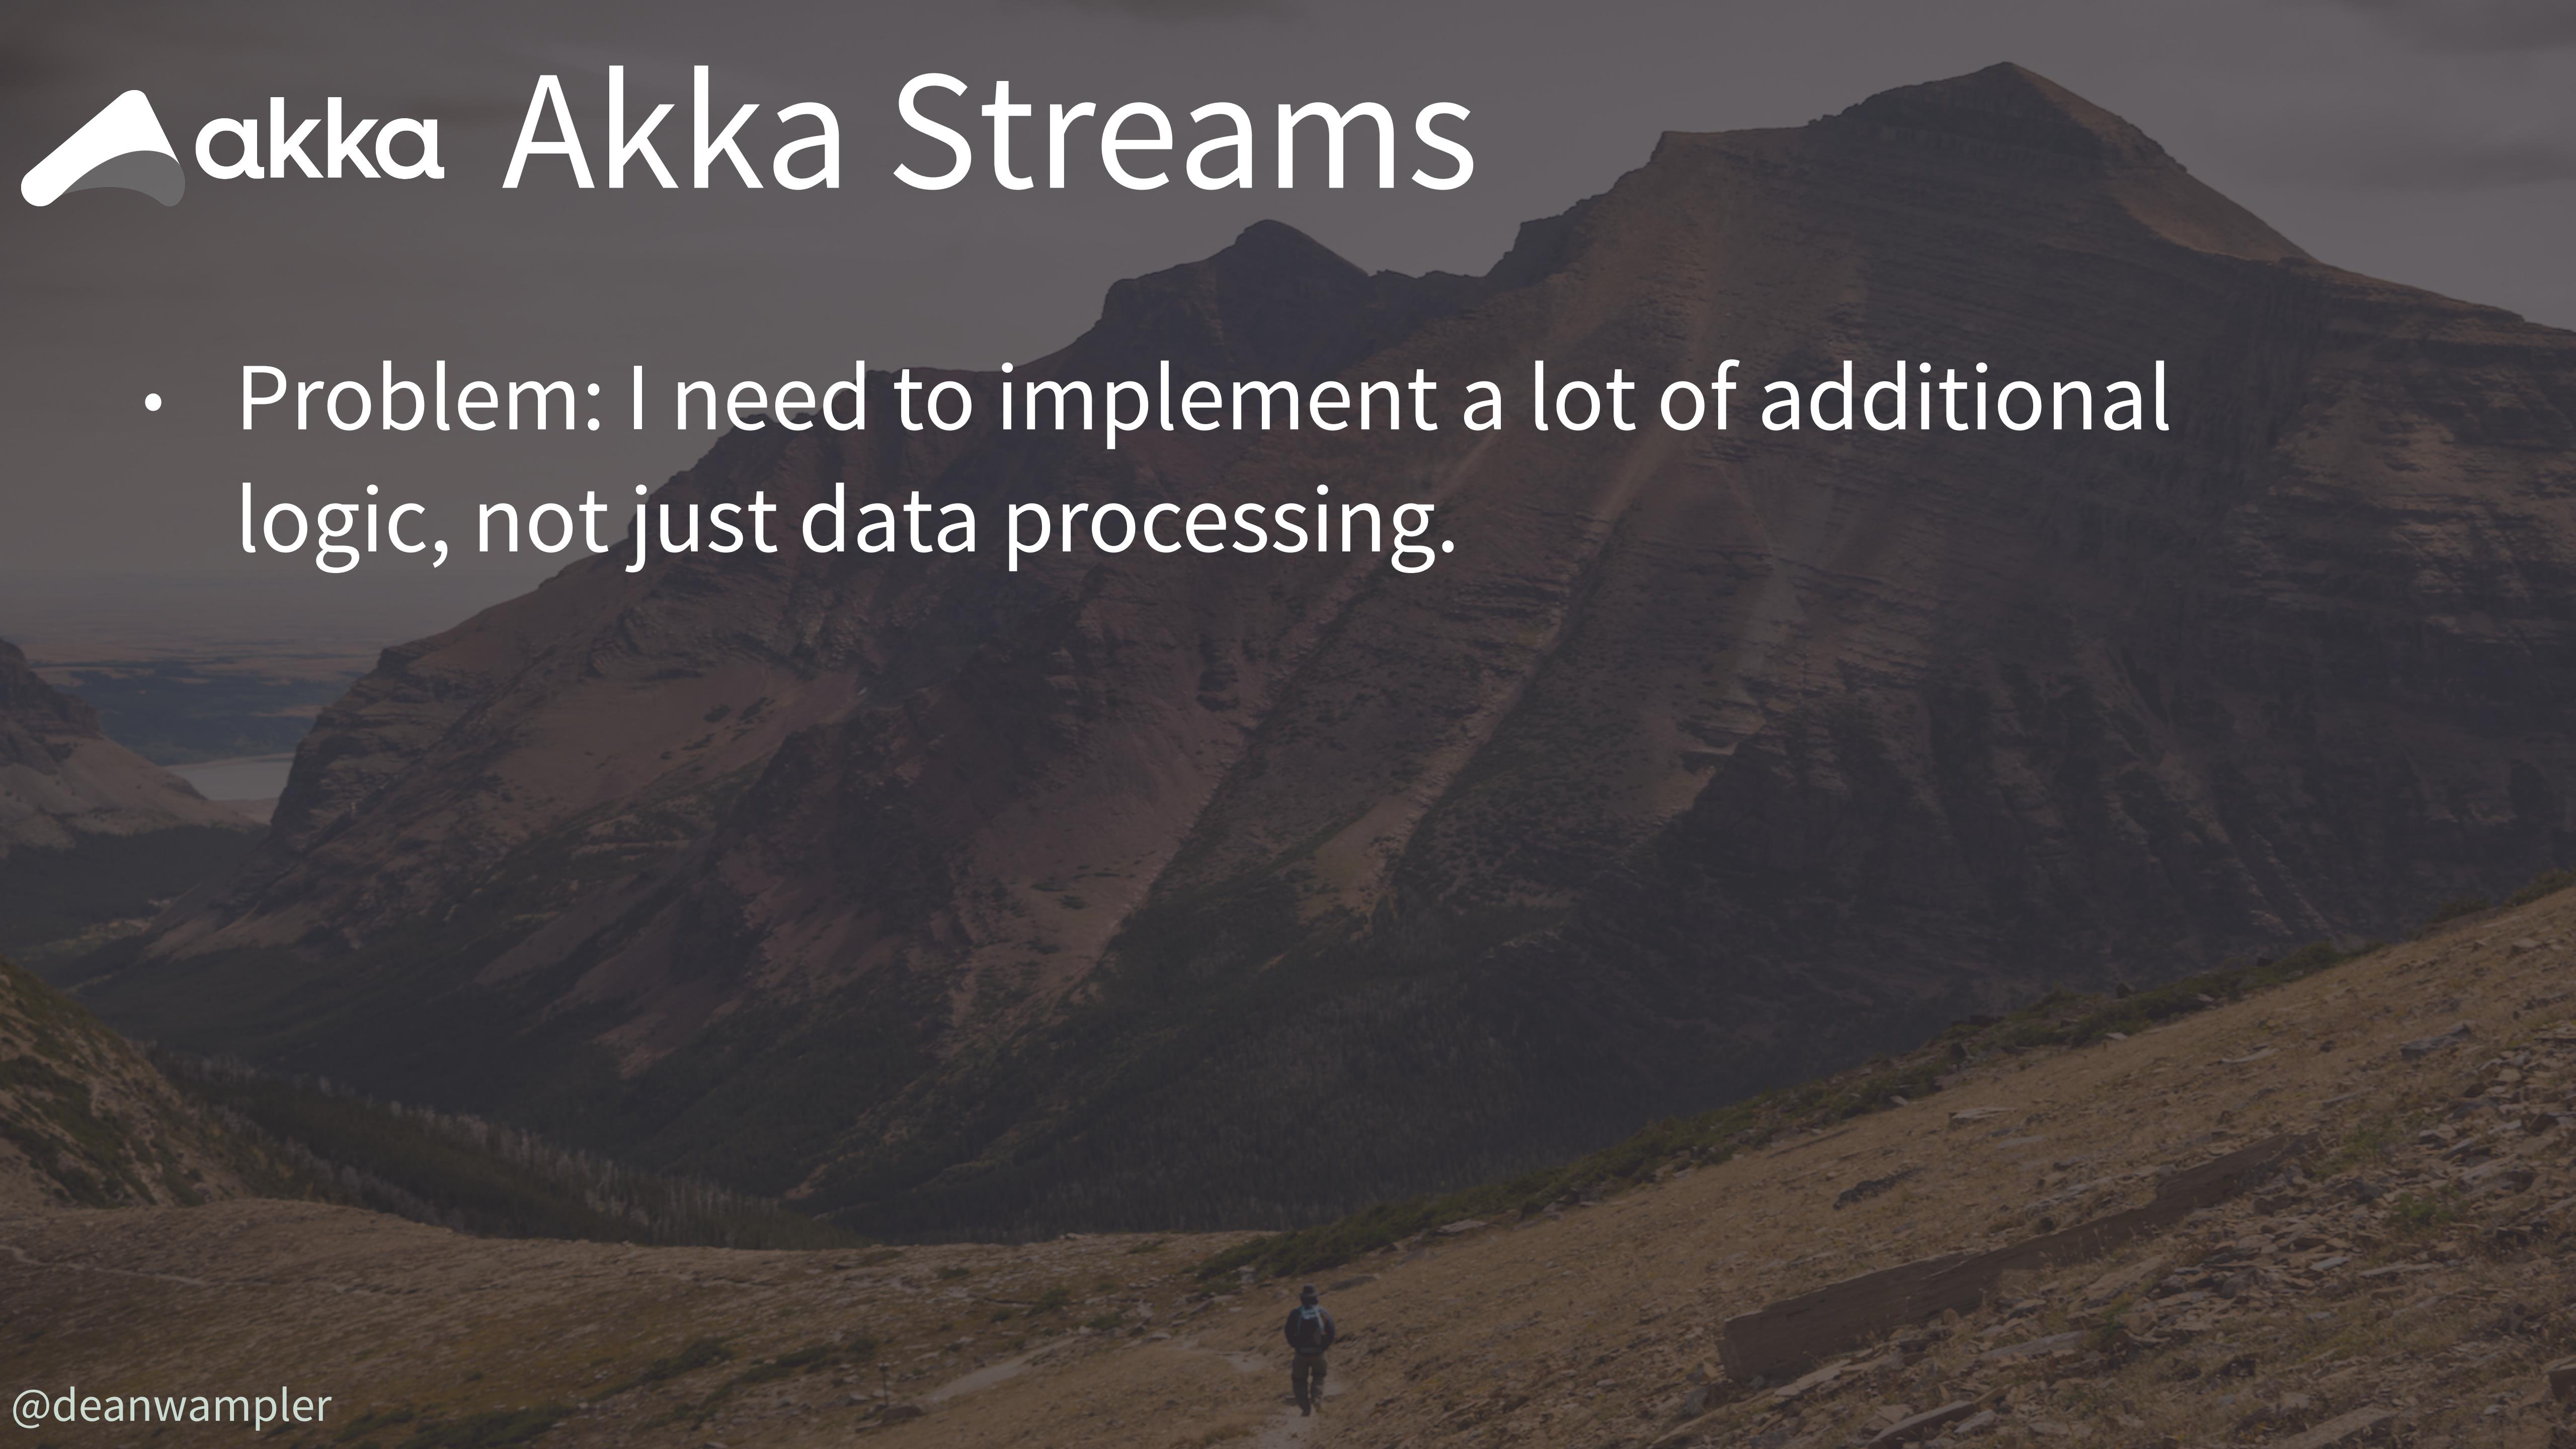
<http://www.reactive-streams.org/>



# Akka Streams



Back pressure composes!



# akka Akka Streams

- Problem: I need to implement a lot of additional logic, not just data processing.



# akka Akka Streams

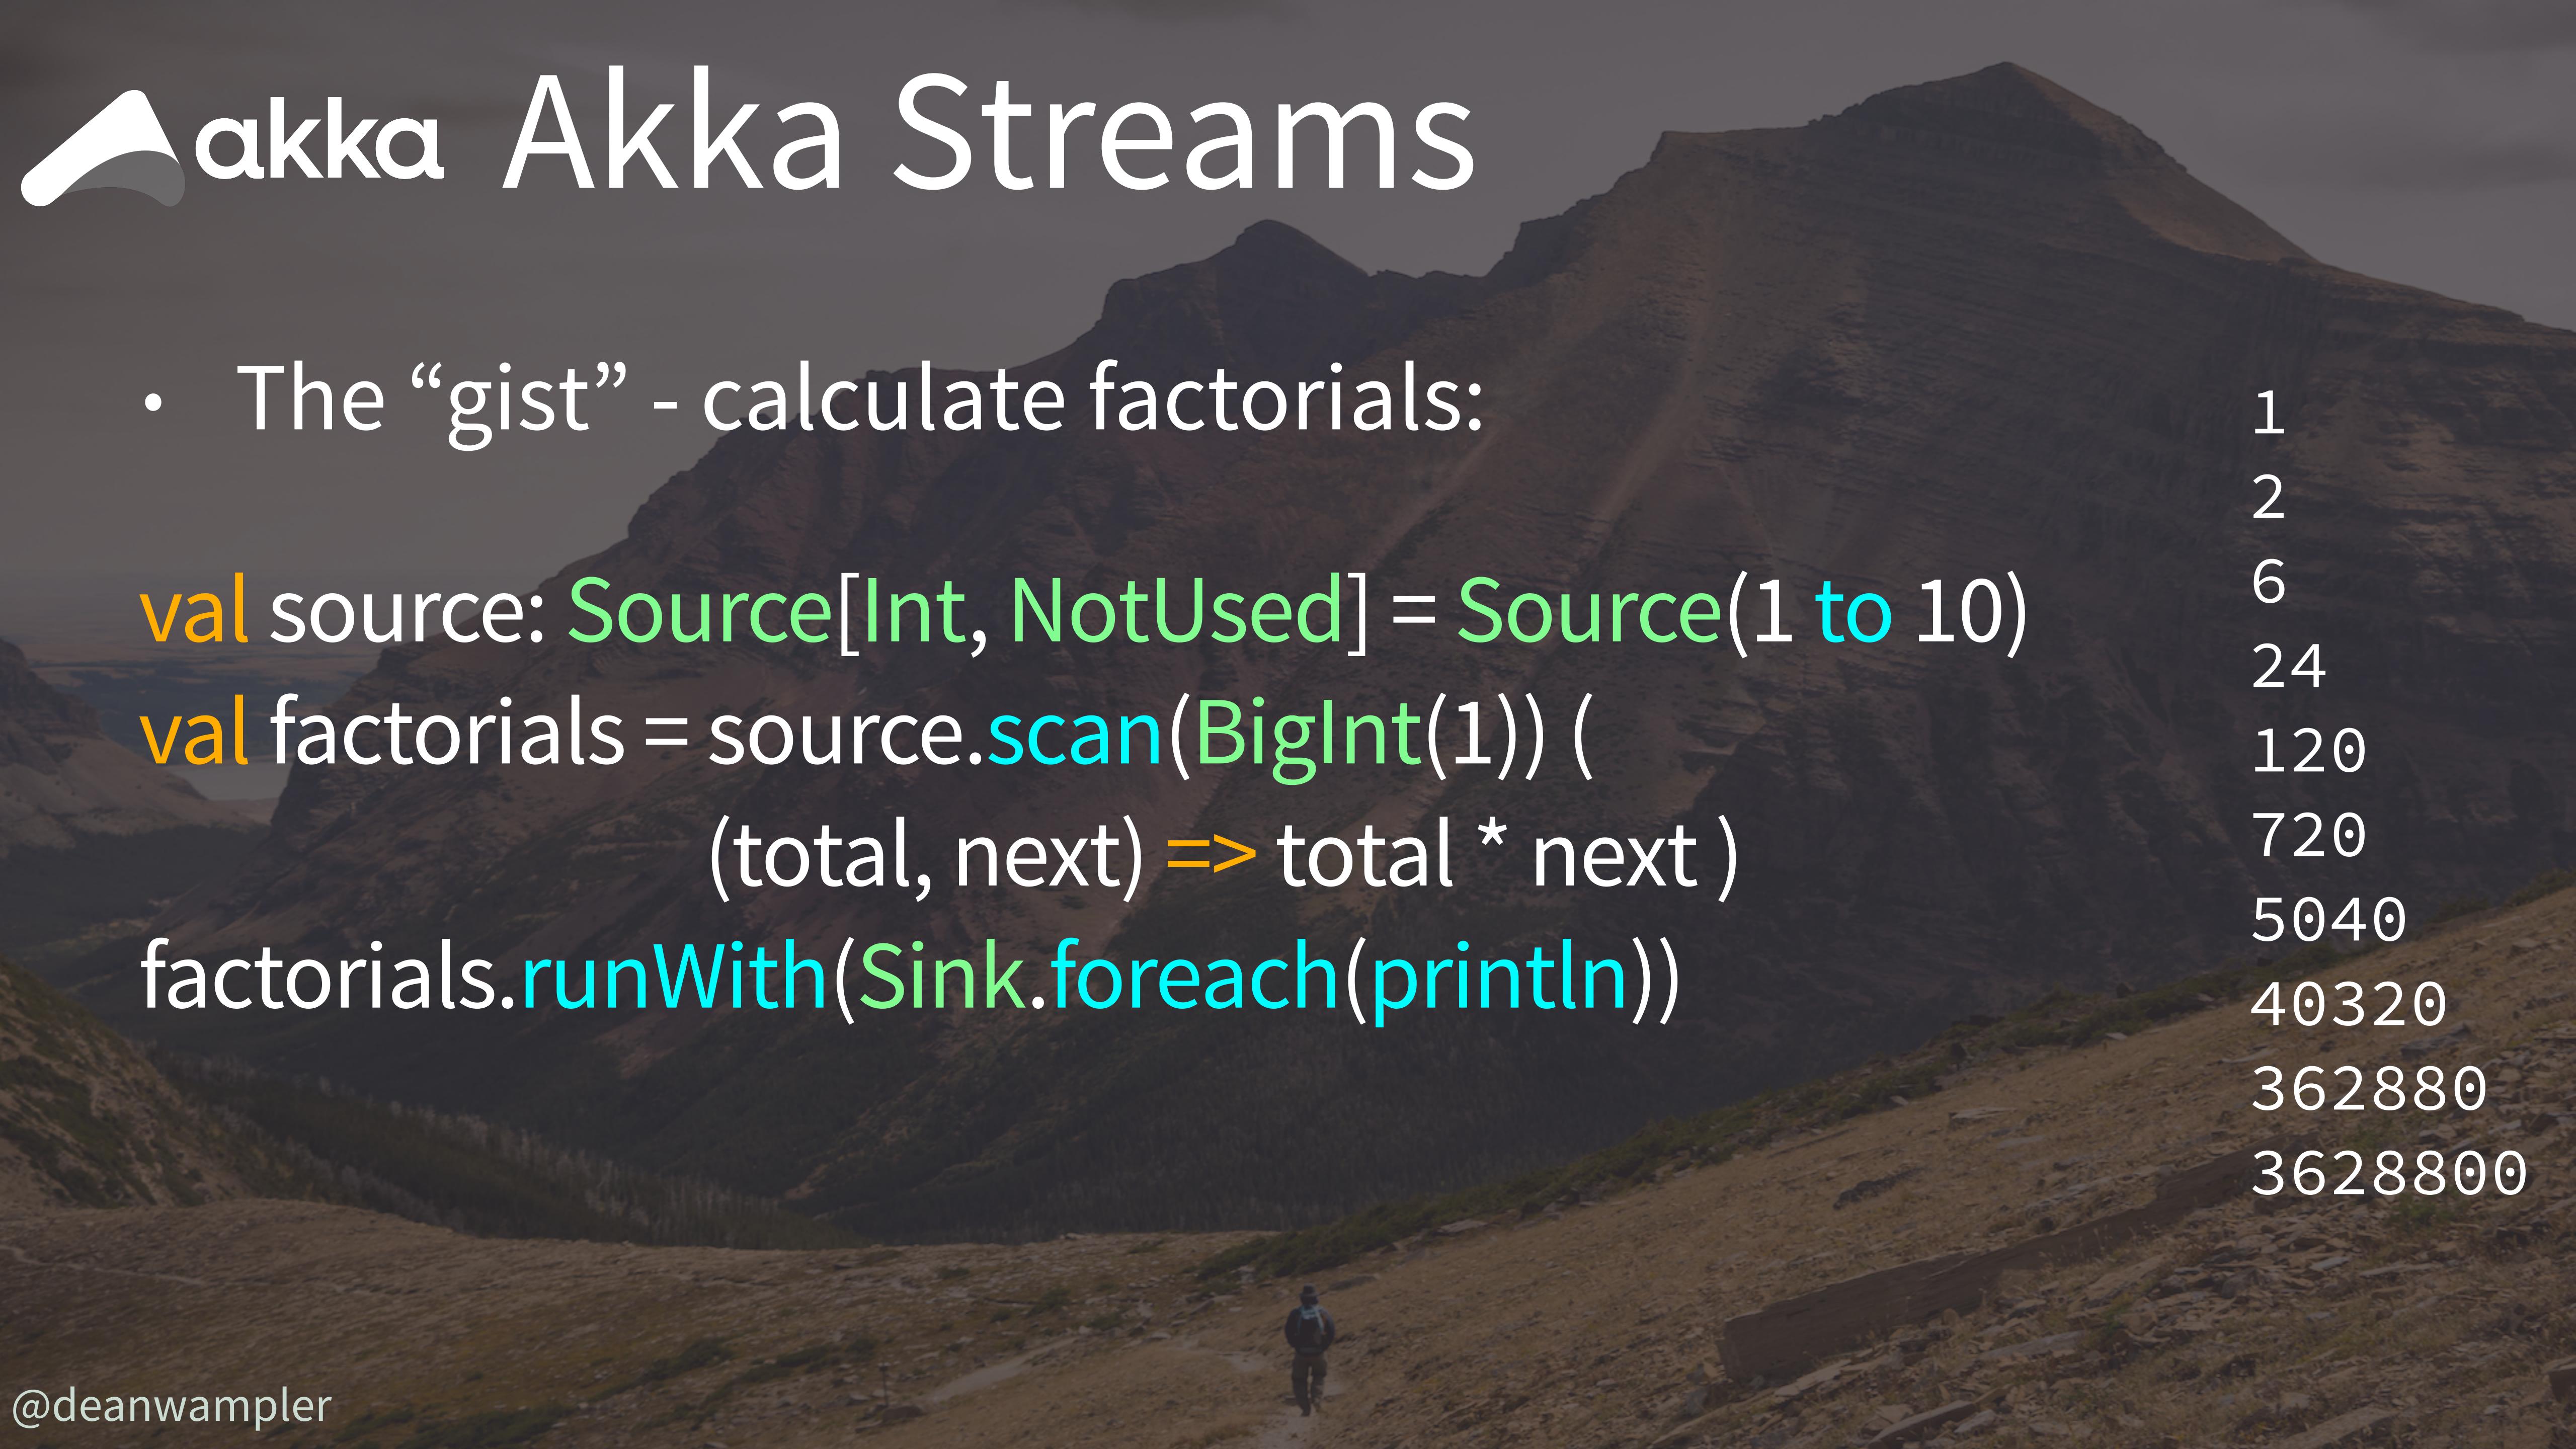
- Solution: Part of the Akka ecosystem
  - Akka Actors, Akka Cluster, Akka HTTP, Akka Persistence, ...
  - Alpakka - rich connection library
  - Optimized for low overhead and latency



# Akka Streams

- The “gist” - calculate factorials:

```
val source: Source[Int, NotUsed] = Source(1 to 10)
val factorials = source.scan(BigInt(1)) (
  (total, next) => total * next )
factorials.runWith(Sink.foreach(println))
```

A scenic mountain landscape with a hiker in the foreground. The mountains are rugged and rocky, with patches of snow and green vegetation. The sky is clear and blue. A small figure of a person is walking on a path in the lower-left foreground.

1

2

6

24

120

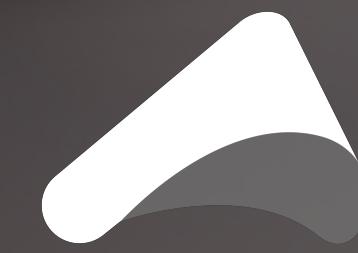
720

5040

40320

362880

3628800



# Akka Streams

- The “gist” - calculate factorials:

```
val source: Source[Int, NotUsed] = Source(1 to 10)
```

```
val factorials = source.scan(BigInt(1)) (  
    (total, next) => total * next )
```

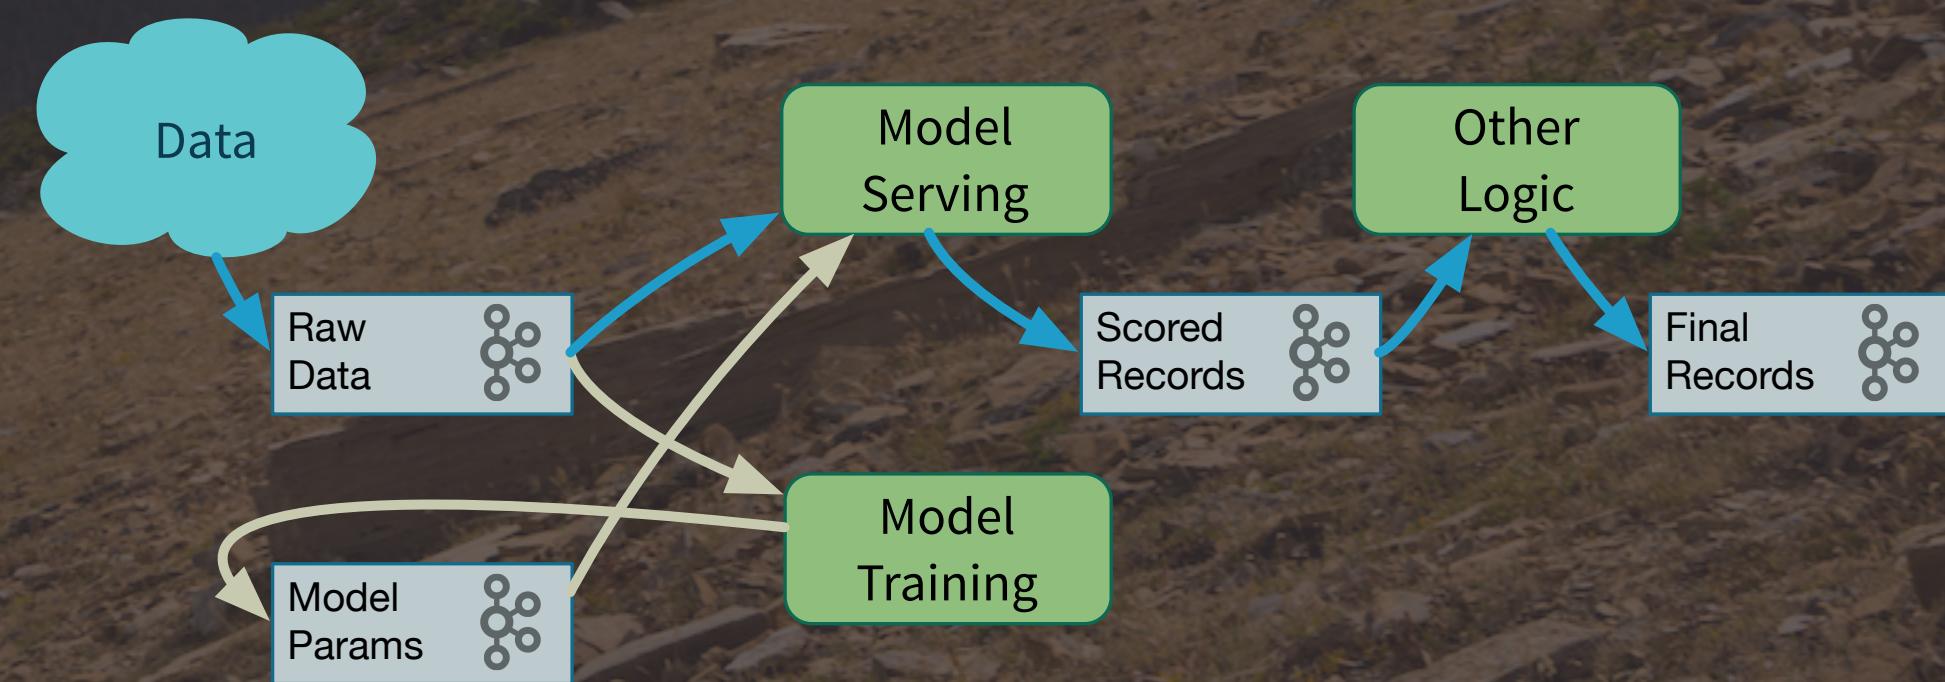
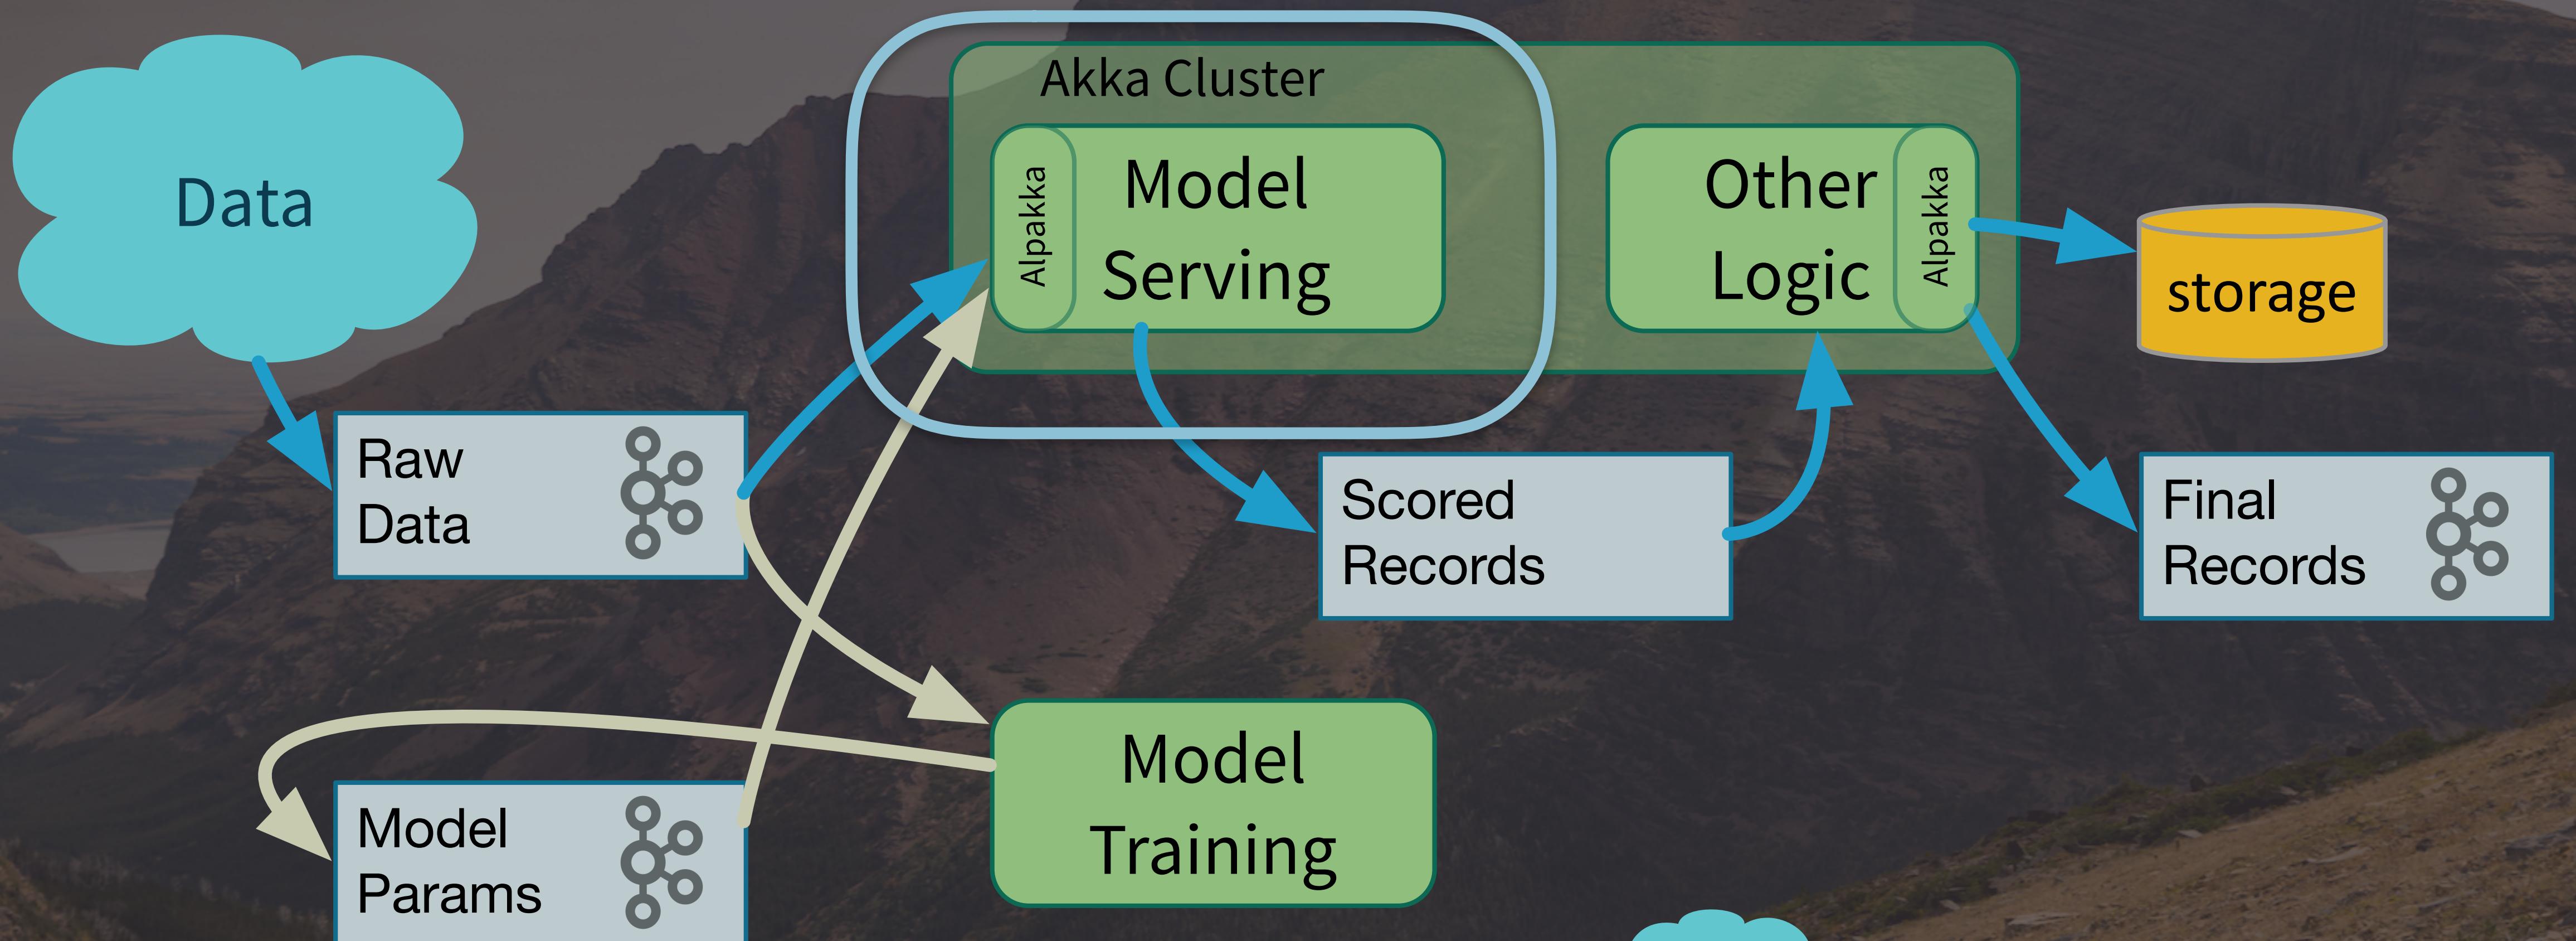
```
factorials.runWith(Sink.foreach(println))
```

A “Graph”

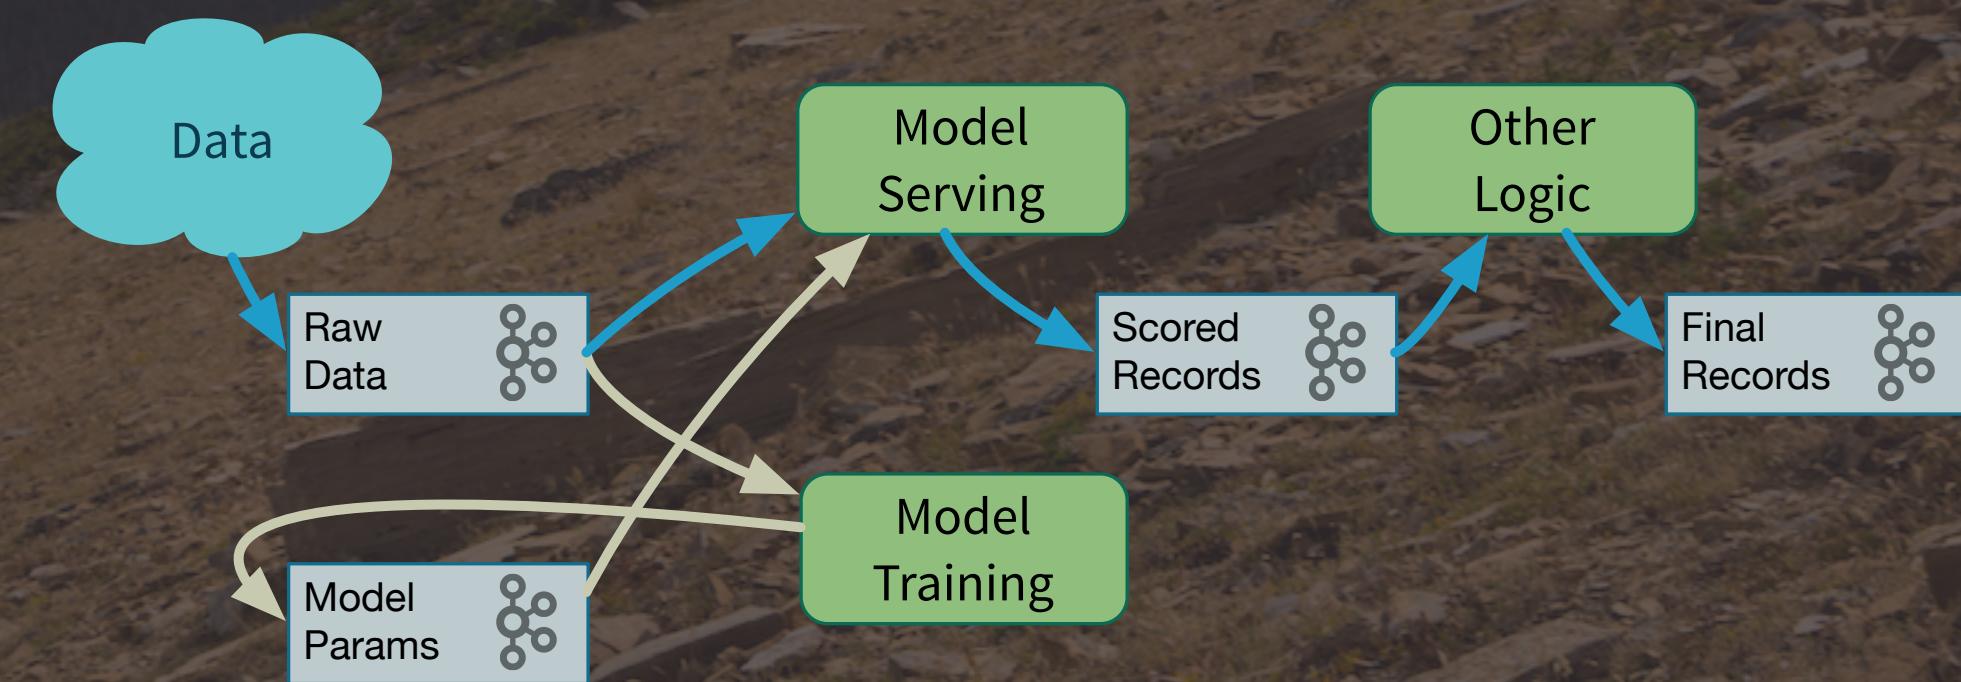
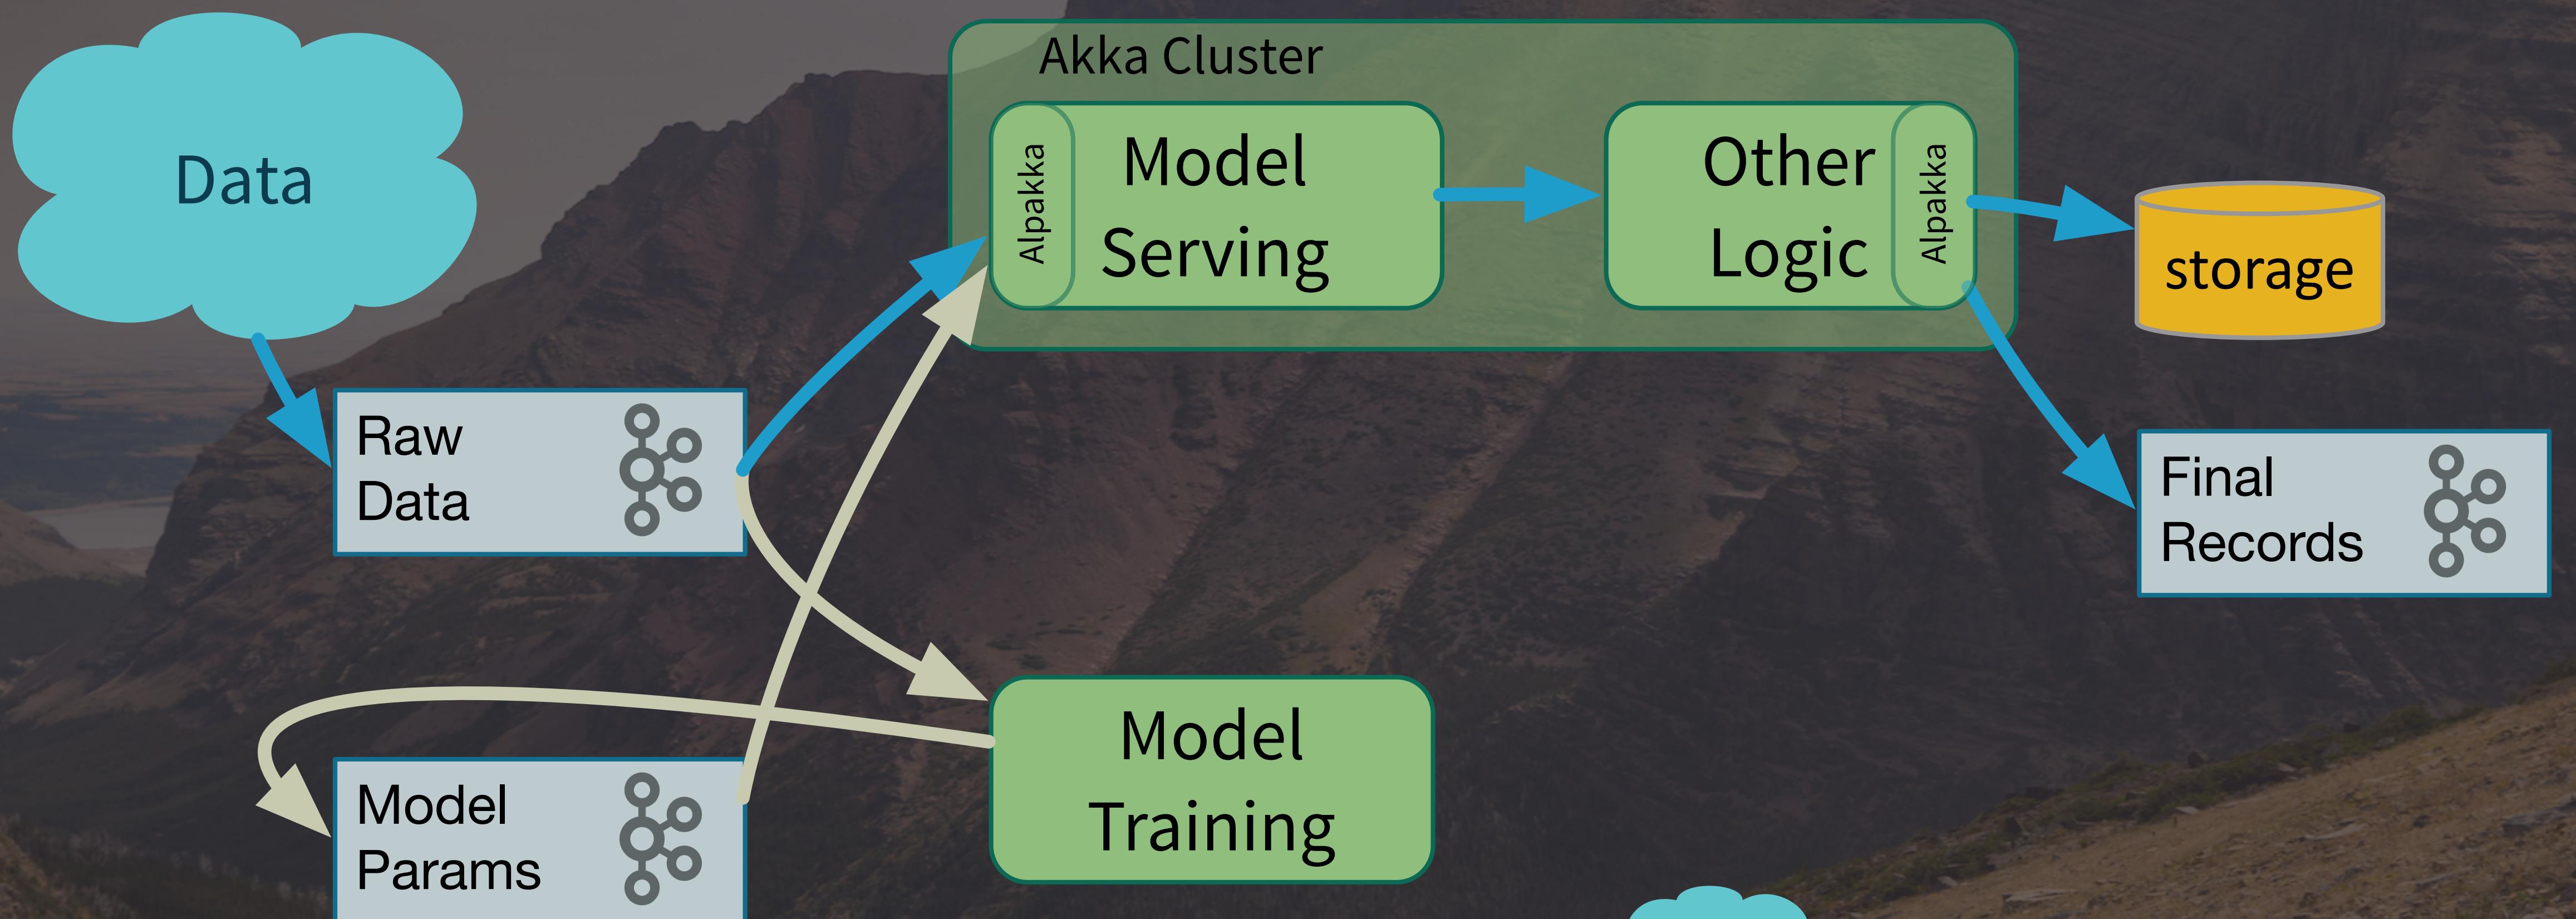


1  
2  
6  
24  
120  
720  
5040  
40320  
362880  
3628800

# Akka Streams Example



# Akka Streams Example - “Direct” Variation

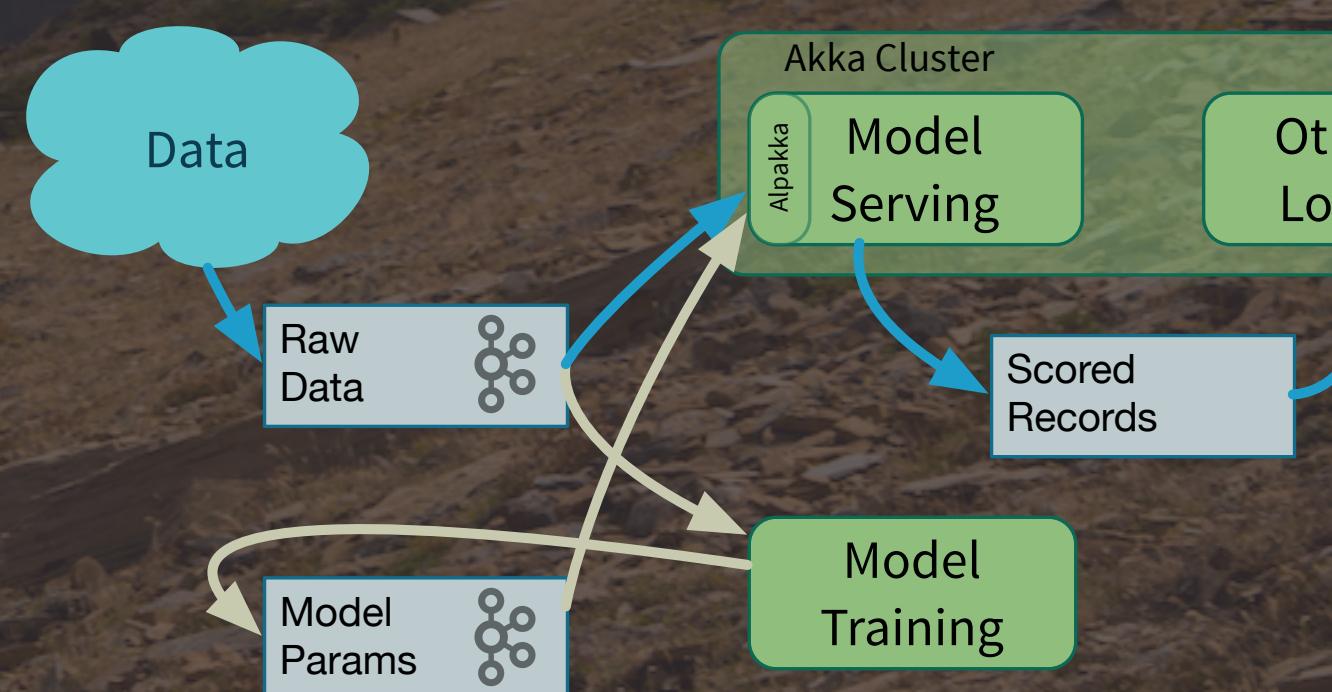


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom “stage”
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

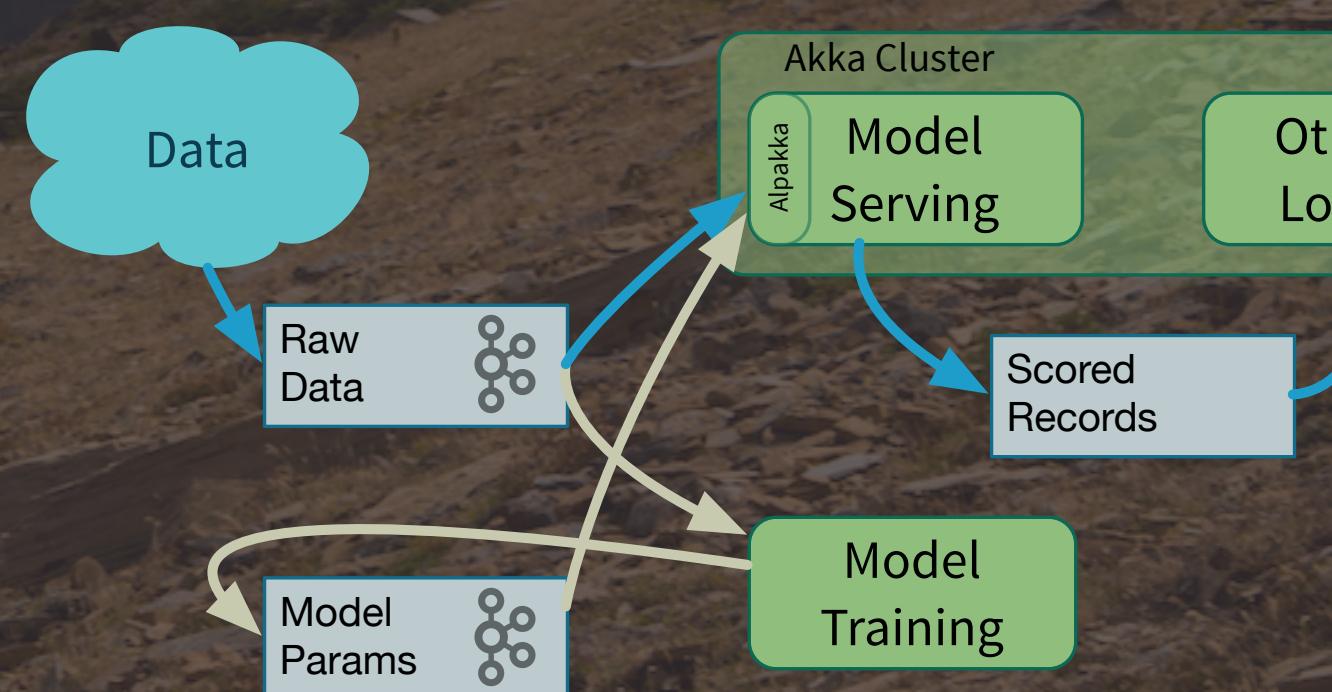


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

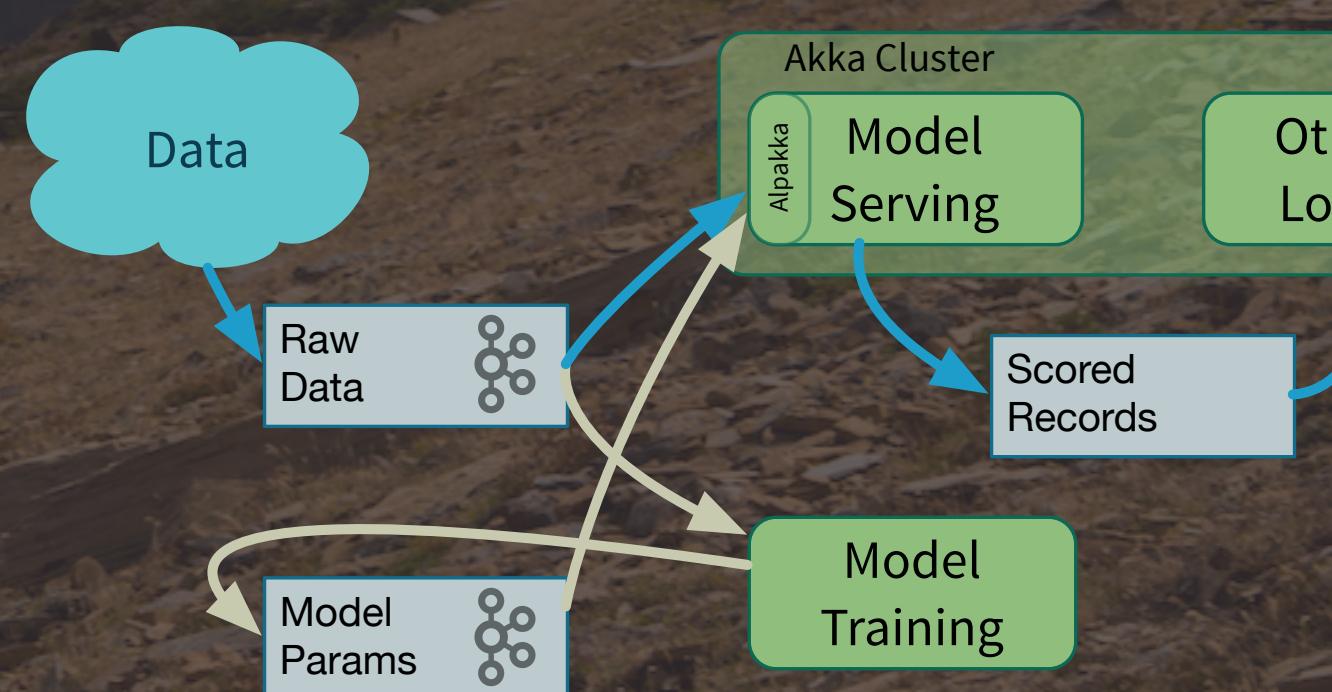


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom "stage"
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

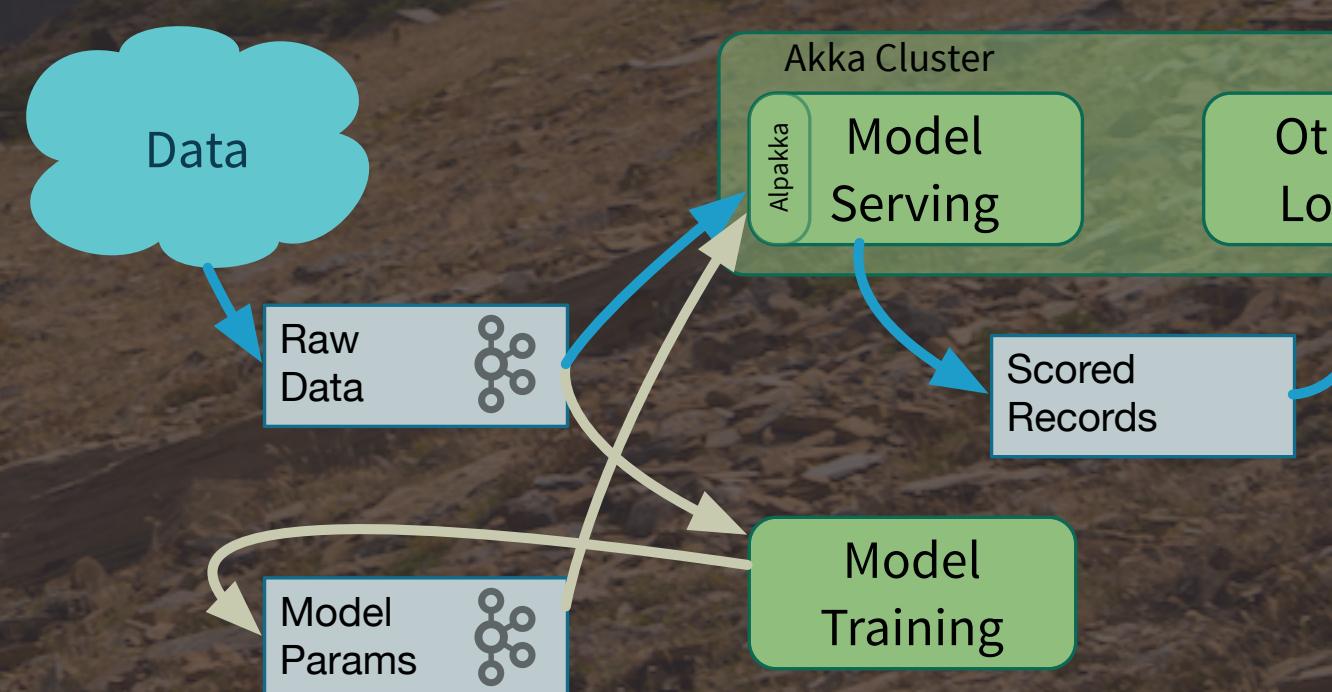


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom “stage”
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```

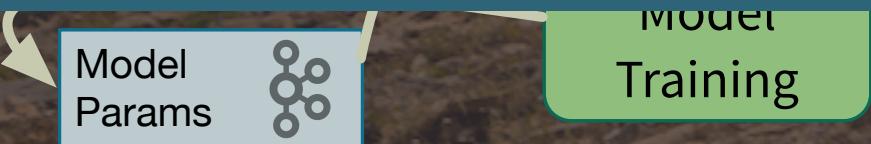


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
import spray.actor.executionContext -> system.dispatcher
case class ModelScoringStage(scorer: ...) extends
  GraphStageWithMaterializedValue[..., ...] {

  val dataRecordIn = Inlet[Record]("dataRecordIn")
  val modelRecordIn = Inlet[ModelImpl]("modelRecordIn")
  val scoringResultOut = Outlet[ScoredRecord]("scoringOut")

  ...
  setHandler(dataRecordIn, new InHandler {
    override def onPush(): Unit = {
      val record = grab(dataRecordIn)
      val newRecord = new ScoredRecord(scorer.score(record), record))
      push(scoringResultOut, Some(newRecord))
      pull(dataRecordIn)
    }
  })
  ...
}

.collect{ case Success(mod) => mod }
.map(model => ModelImpl.findModel(model))
.collect{ case Success(modImpl) => modImpl }
```

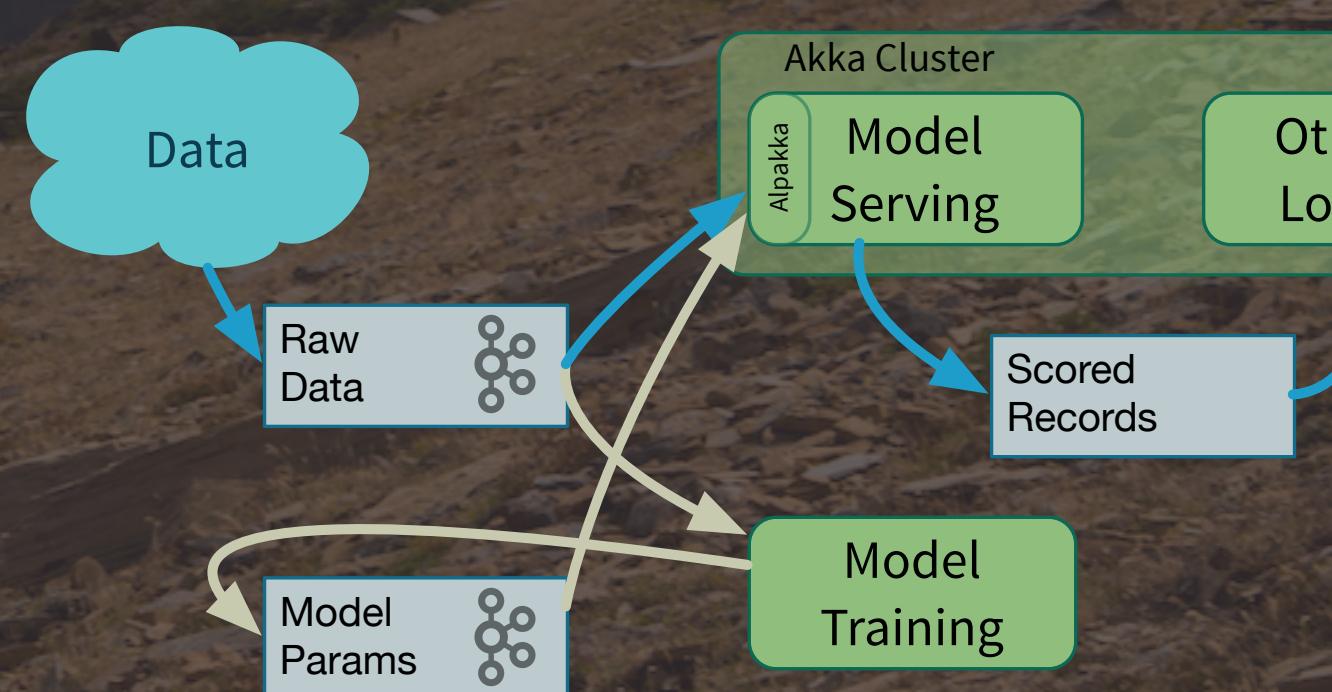


```
implicit val system = ActorSystem("ModelServing")
implicit val materializer = ActorMaterializer()
implicit val executionContext = system.dispatcher
```

```
val modelProcessor = new ModelProcessor // Same as KS example
val scorer = new Scorer(modelProcessor) // Same as KS example
val modelScoringStage = new ModelScoringStage(scorer)// AS custom “stage”
```

```
val dataStream: Source[Record, Consumer.Control] =
  Consumer.atMostOnceSource(dataConsumerSettings,
    Subscriptions.topics(rawDataTopic))
  .map(input => DataRecord.parseBytes(input.value()))
  .collect{ case Success(data) => data }
```

```
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
```



```

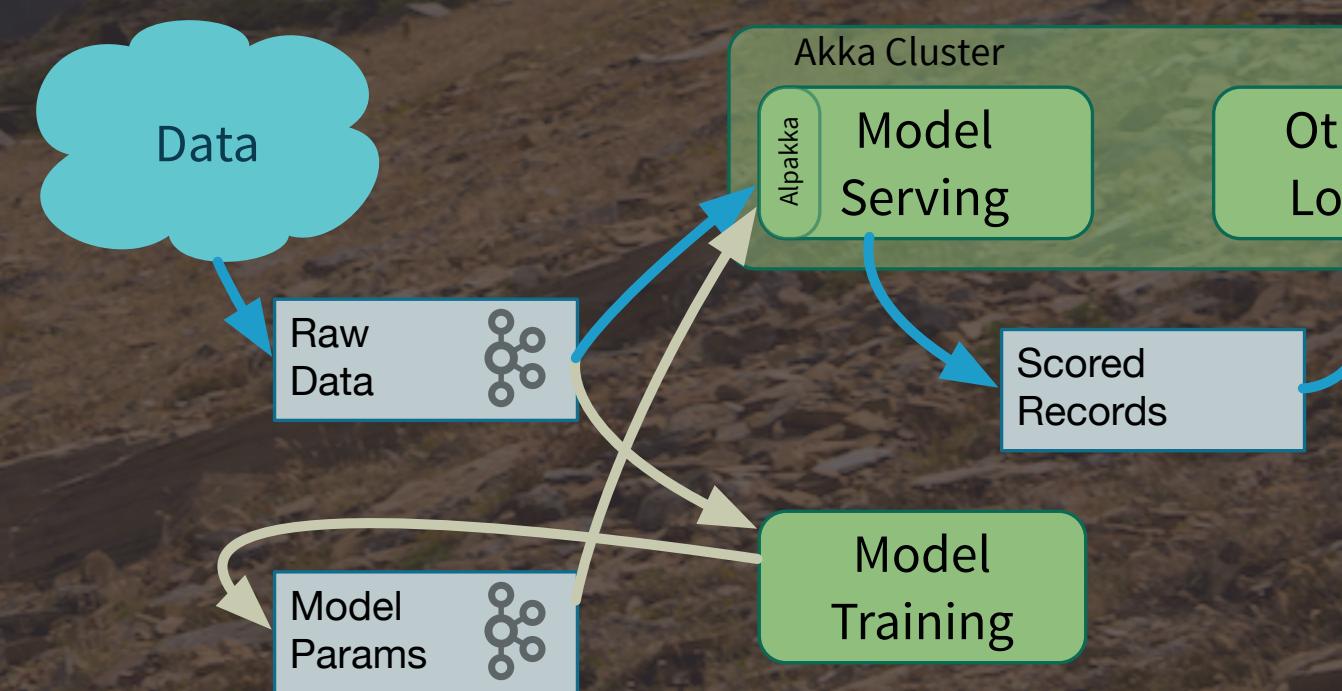
val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
  .foreach(modImpl => modelProcessor.setModel(modImpl))
modelStream.to(Sink.ignore).run() // No “sinking” required; just run

```

```

dataStream
  .viaMat(modelScoringStage)(Keep.right)
  .map(result => new ProducerRecord[Array[Byte], ScoredRecord](
    scoredRecordsTopic, result))
  .runWith(Producer.plainSink(producerSettings))

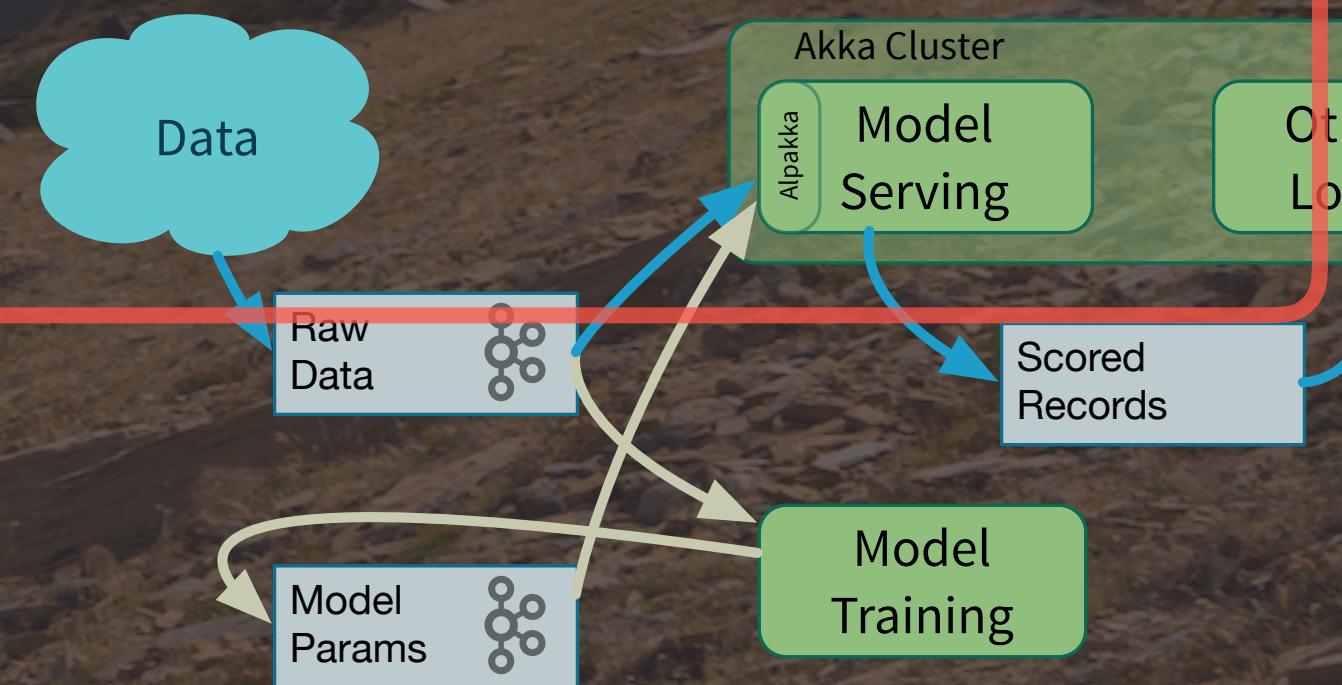
```



```
.collect{ case Success(data) -> data }

val modelStream: Source[ModelImpl, Consumer.Control] =
  Consumer.atMostOnceSource(modelConsumerSettings,
    Subscriptions.topics(modelTopic))
  .map(input => Model.parseBytes(input.value()))
  .collect{ case Success(mod) => mod }
  .map(model => ModelImpl.findModel(model))
  .collect{ case Success(modImpl) => modImpl }
  .foreach(modImpl => modelProcessor.setModel(modImpl))
modelStream.to(Sink.ignore).run() // No “sinking” required; just run
```

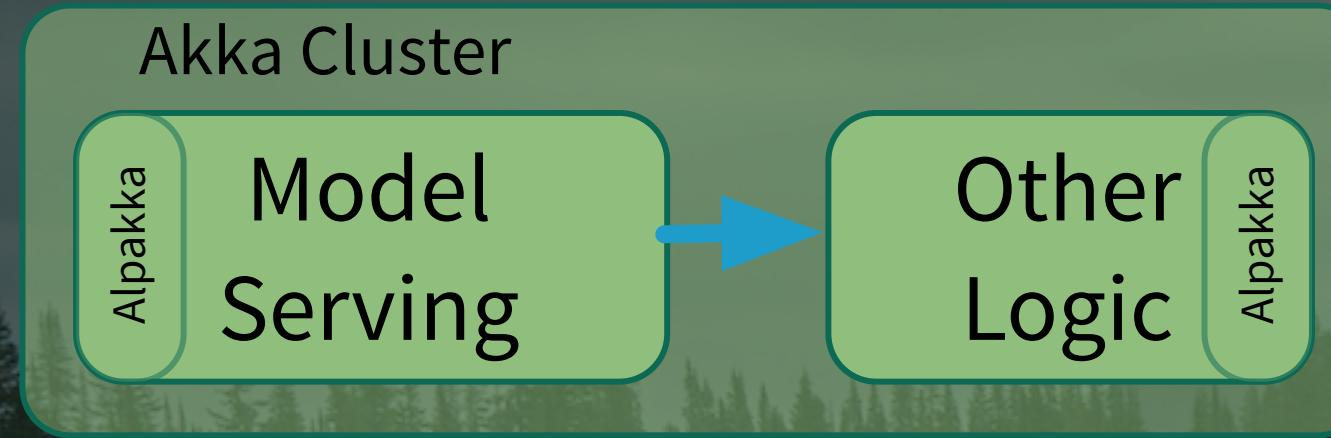
```
dataStream
  .viaMat(modelScoringStage)(Keep.right)
  .map(result => new ProducerRecord[Array[Byte], ScoredRecord](
    scoredRecordsTopic, result))
  .runWith(Producer.plainSink(producerSettings))
```



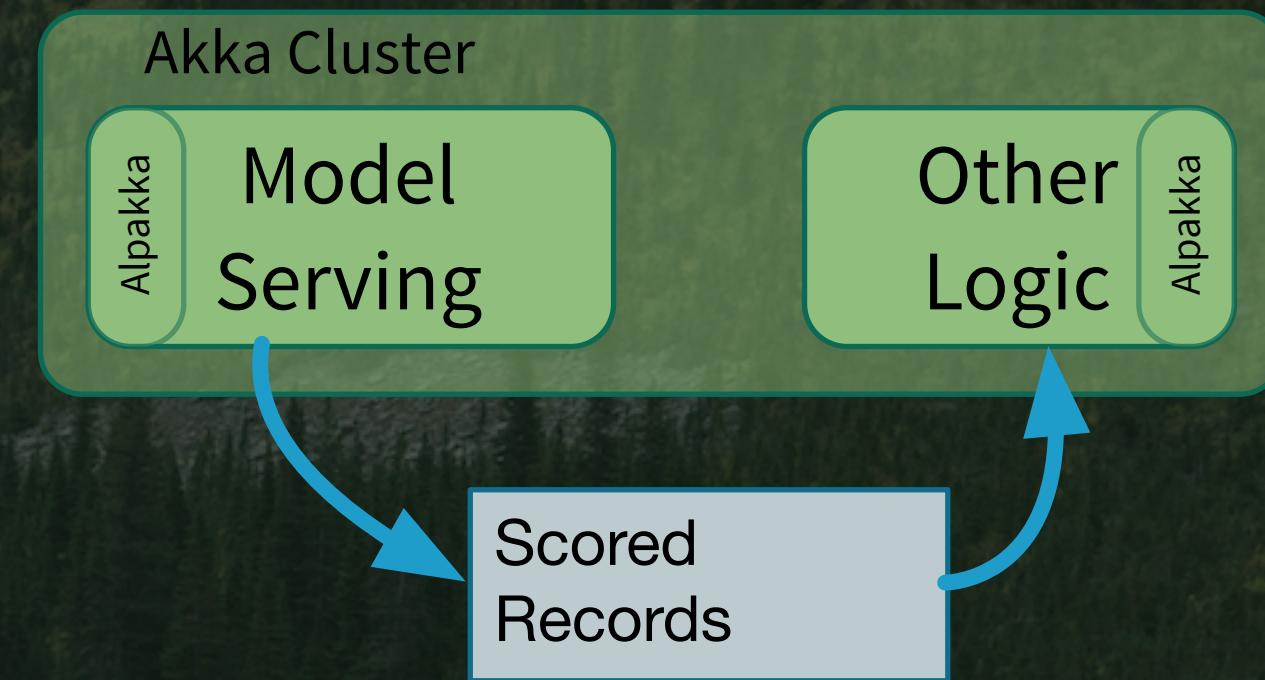
A wide-angle photograph of a mountain lake at sunset. The sky is filled with dramatic, layered clouds. The calm water of the lake perfectly reflects the surrounding landscape, including a dense forest of tall evergreen trees on the left bank and a steep, rocky mountain slope covered in green vegetation on the right. The overall atmosphere is serene and natural.

# Other Concerns

# Go Direct or Through Kafka?



VS.



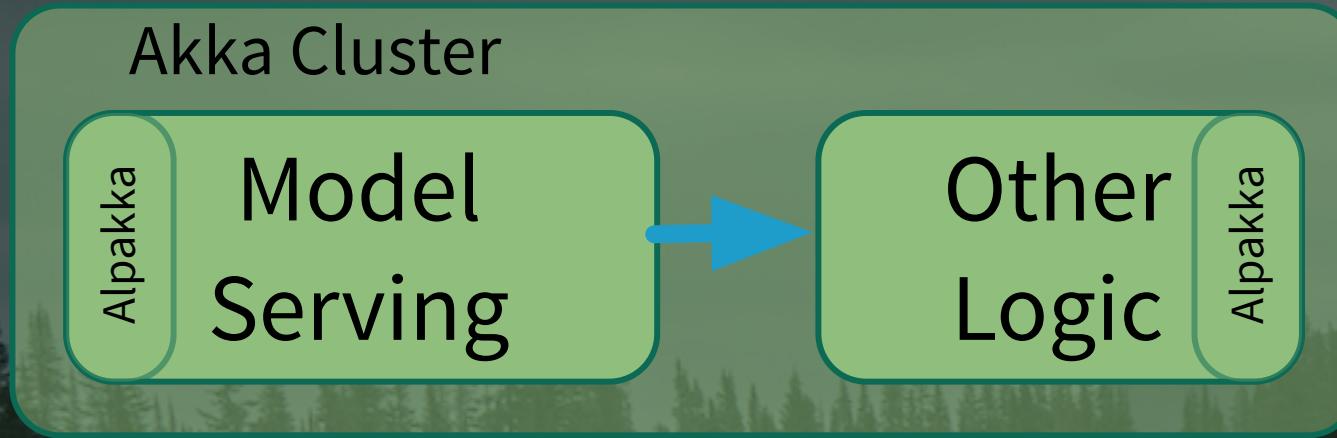
- Extremely low latency
- Higher latency (network, queue depth)

# Go Direct or Through Kafka?

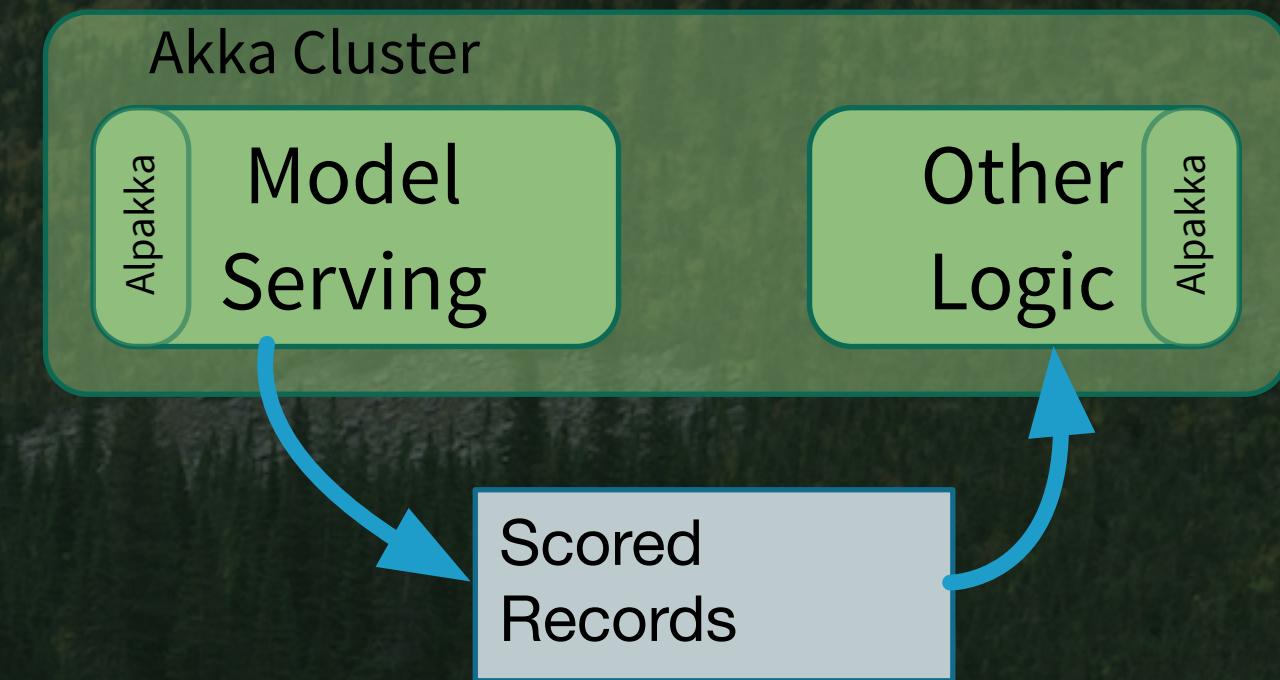


- Extremely low latency
- Minimal I/O and memory overhead. No marshaling overhead
- Higher latency (network, queue depth)
- Higher I/O and processing (marshaling) overhead

# Go Direct or Through Kafka?

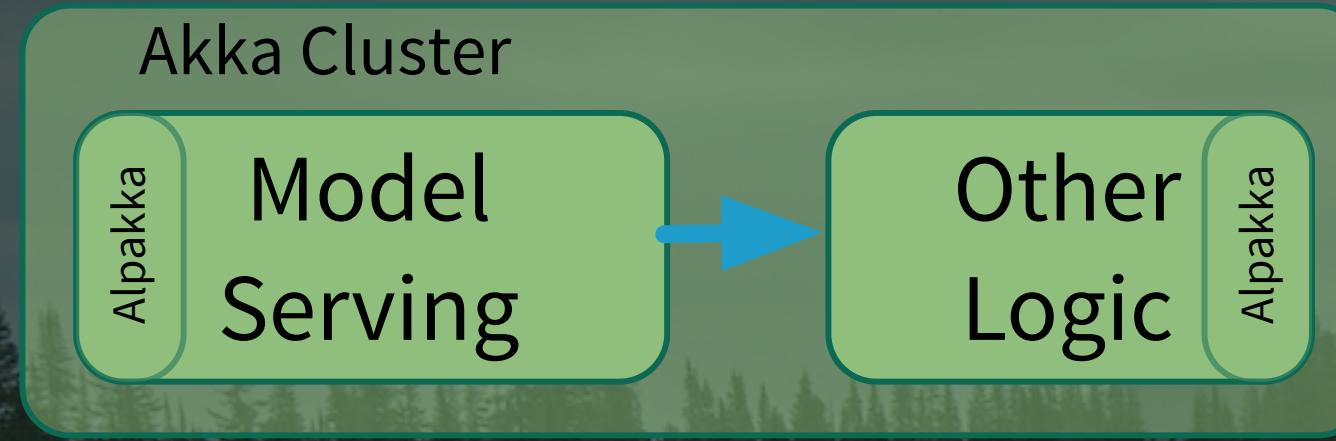


VS.

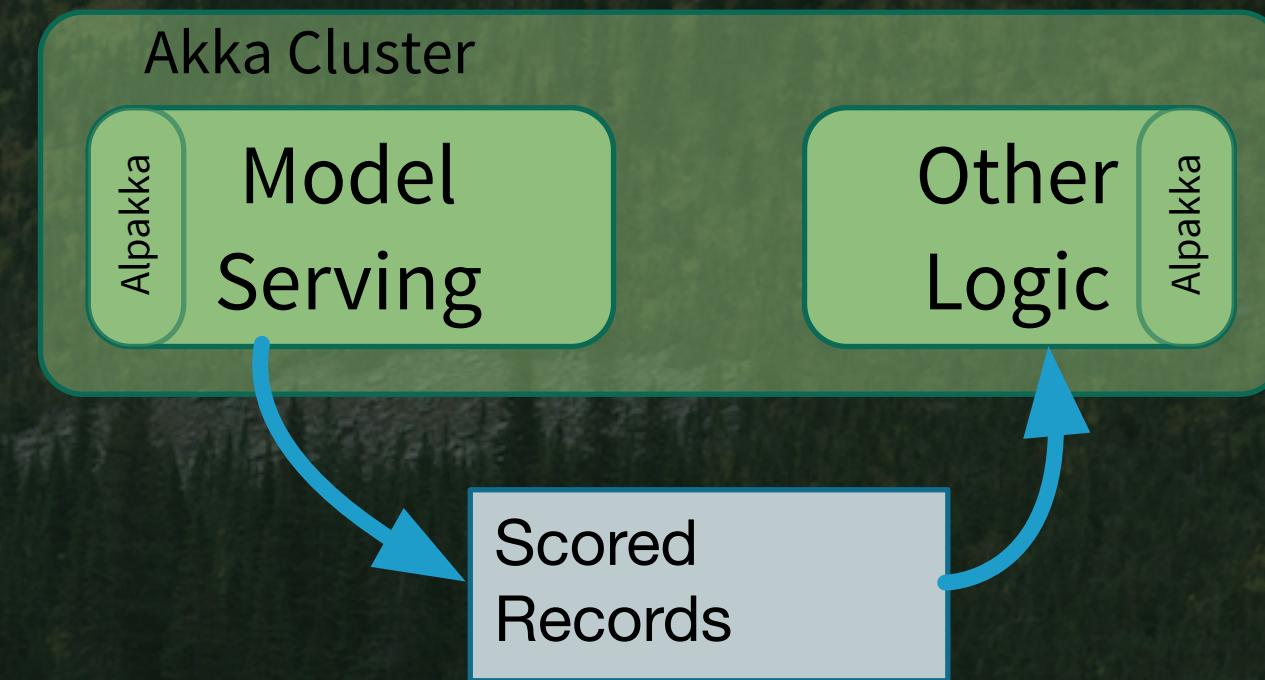


- Extremely low latency
- Minimal I/O and memory overhead. No marshaling overhead
- Hard to scale, evolve independently
- Higher latency (network, queue depth)
- Higher I/O and processing (marshaling) overhead
- Easy independent scalability, evolution

# Go Direct or Through Kafka?



VS.



- *Reactive Streams* back pressure
- Very deep buffer (partition limited by disk size)

# Go Direct or Through Kafka?



VS.

- *Reactive Streams* back pressure
- “Direct” coupling between sender and receiver, but actually uses a URL abstraction
- Very deep buffer (partition limited by disk size)
- Strong decoupling - M producers, N consumers, completely disconnected

?

# Wrapping Up...



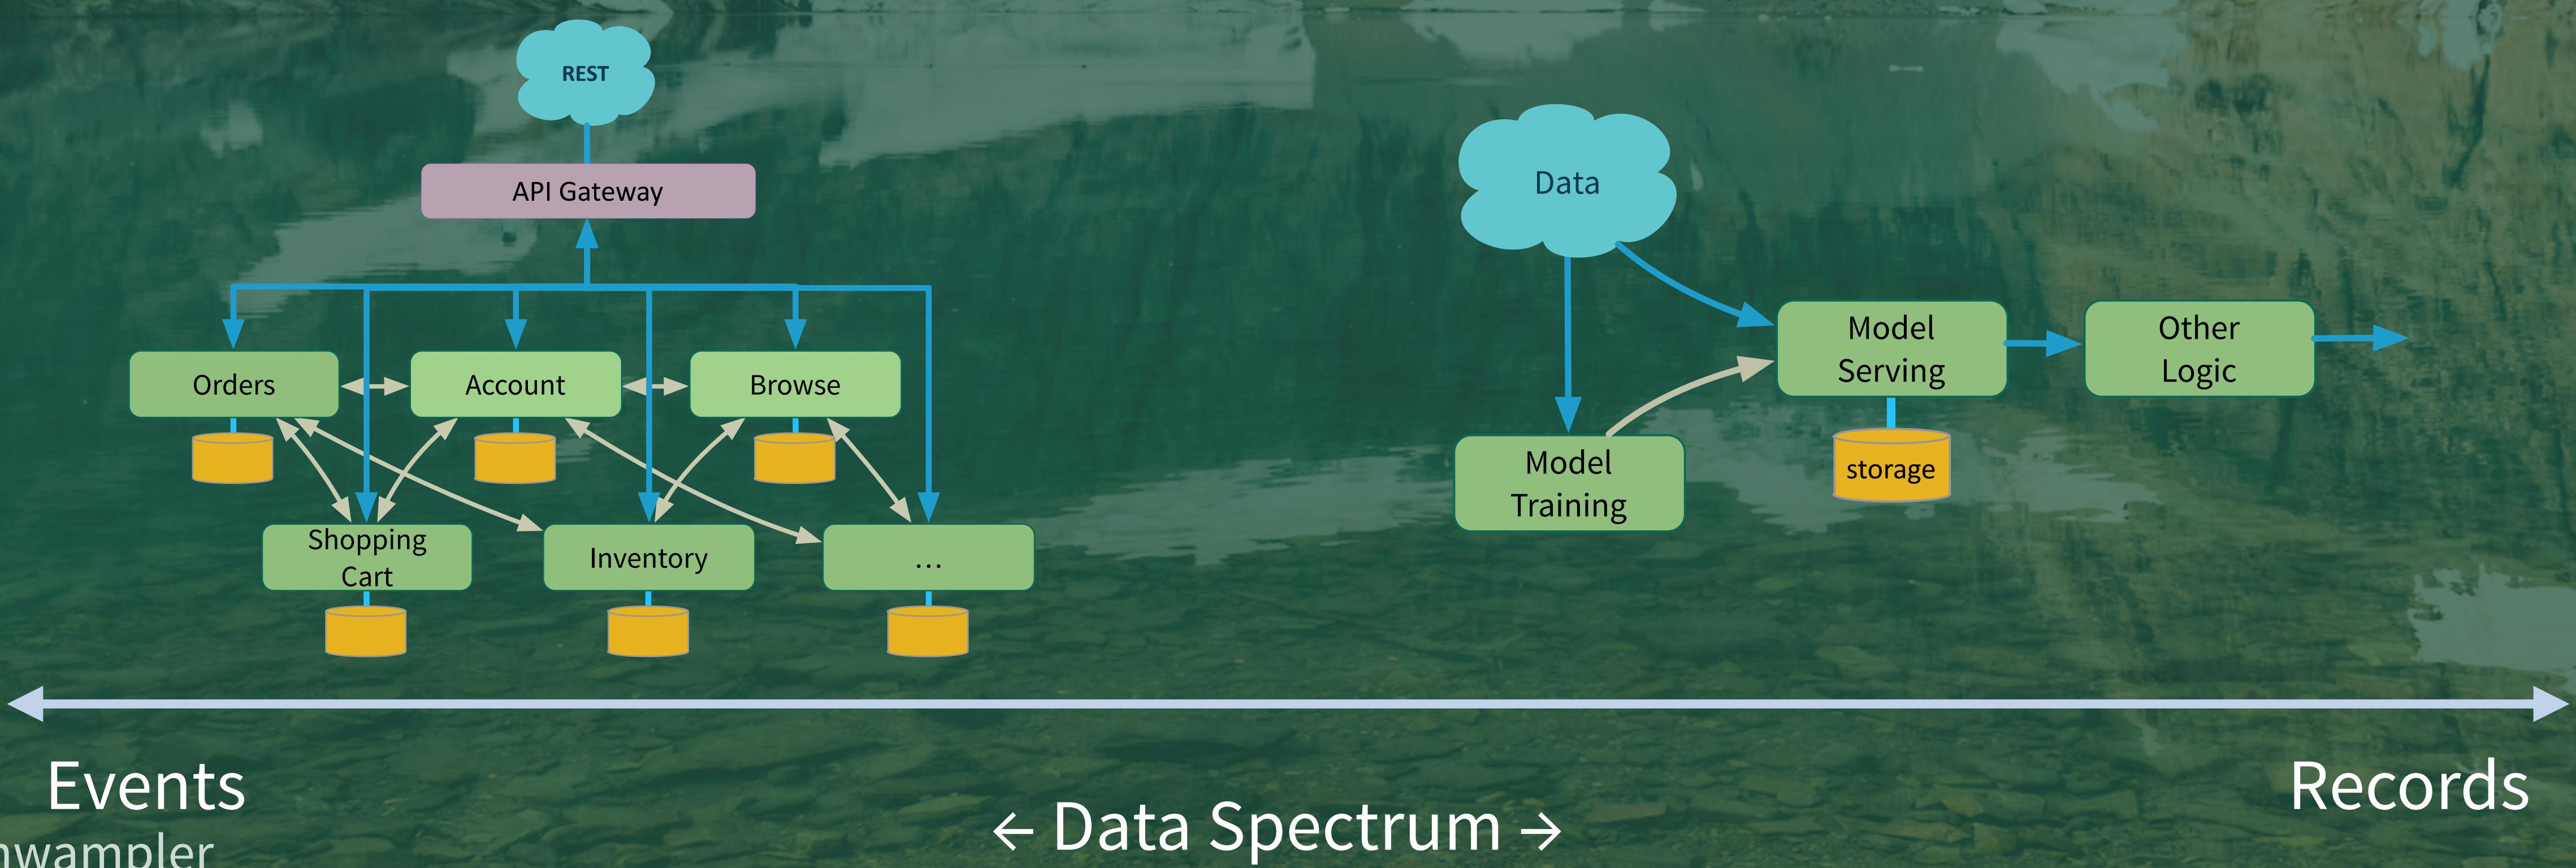
# A Spectrum of Microservices

akka Streams

Streams

Event-driven μ-services

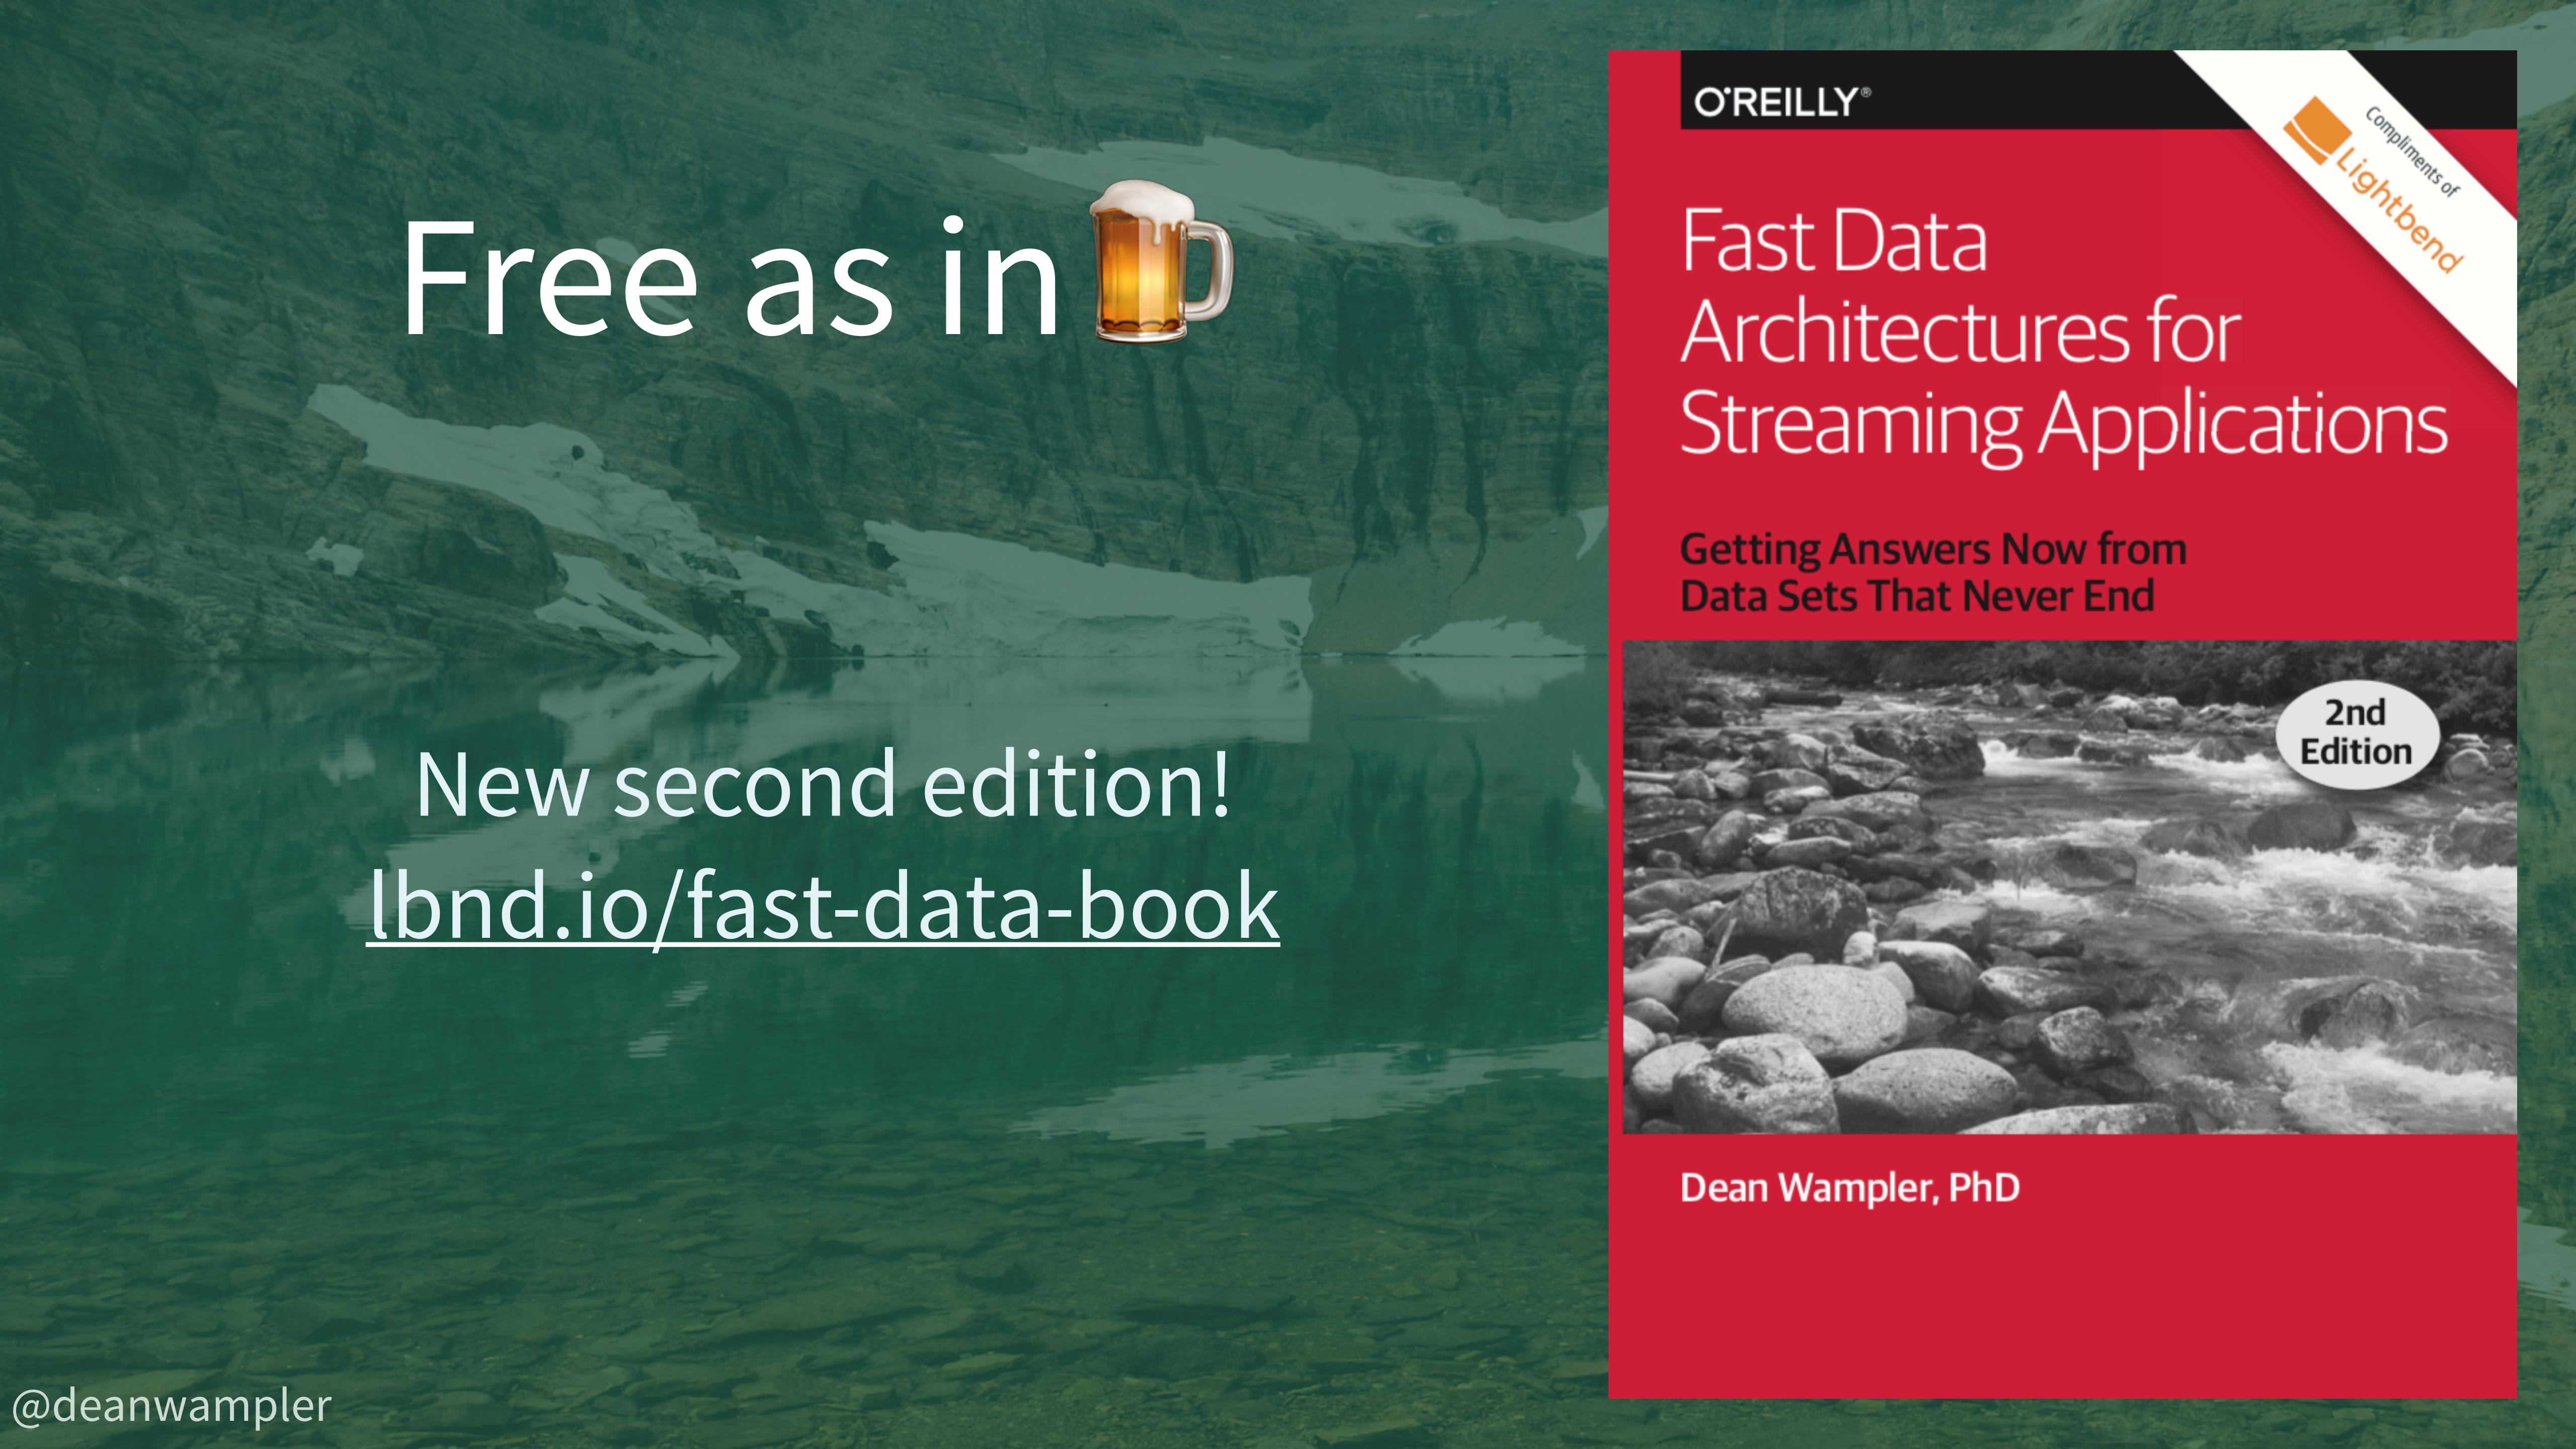
“Record-centric” μ-services



Events

@deanwampler

Records



O'REILLY®

Free as in



New second edition!  
[lbnd.io/fast-data-book](https://lbnd.io/fast-data-book)

# Fast Data Architectures for Streaming Applications

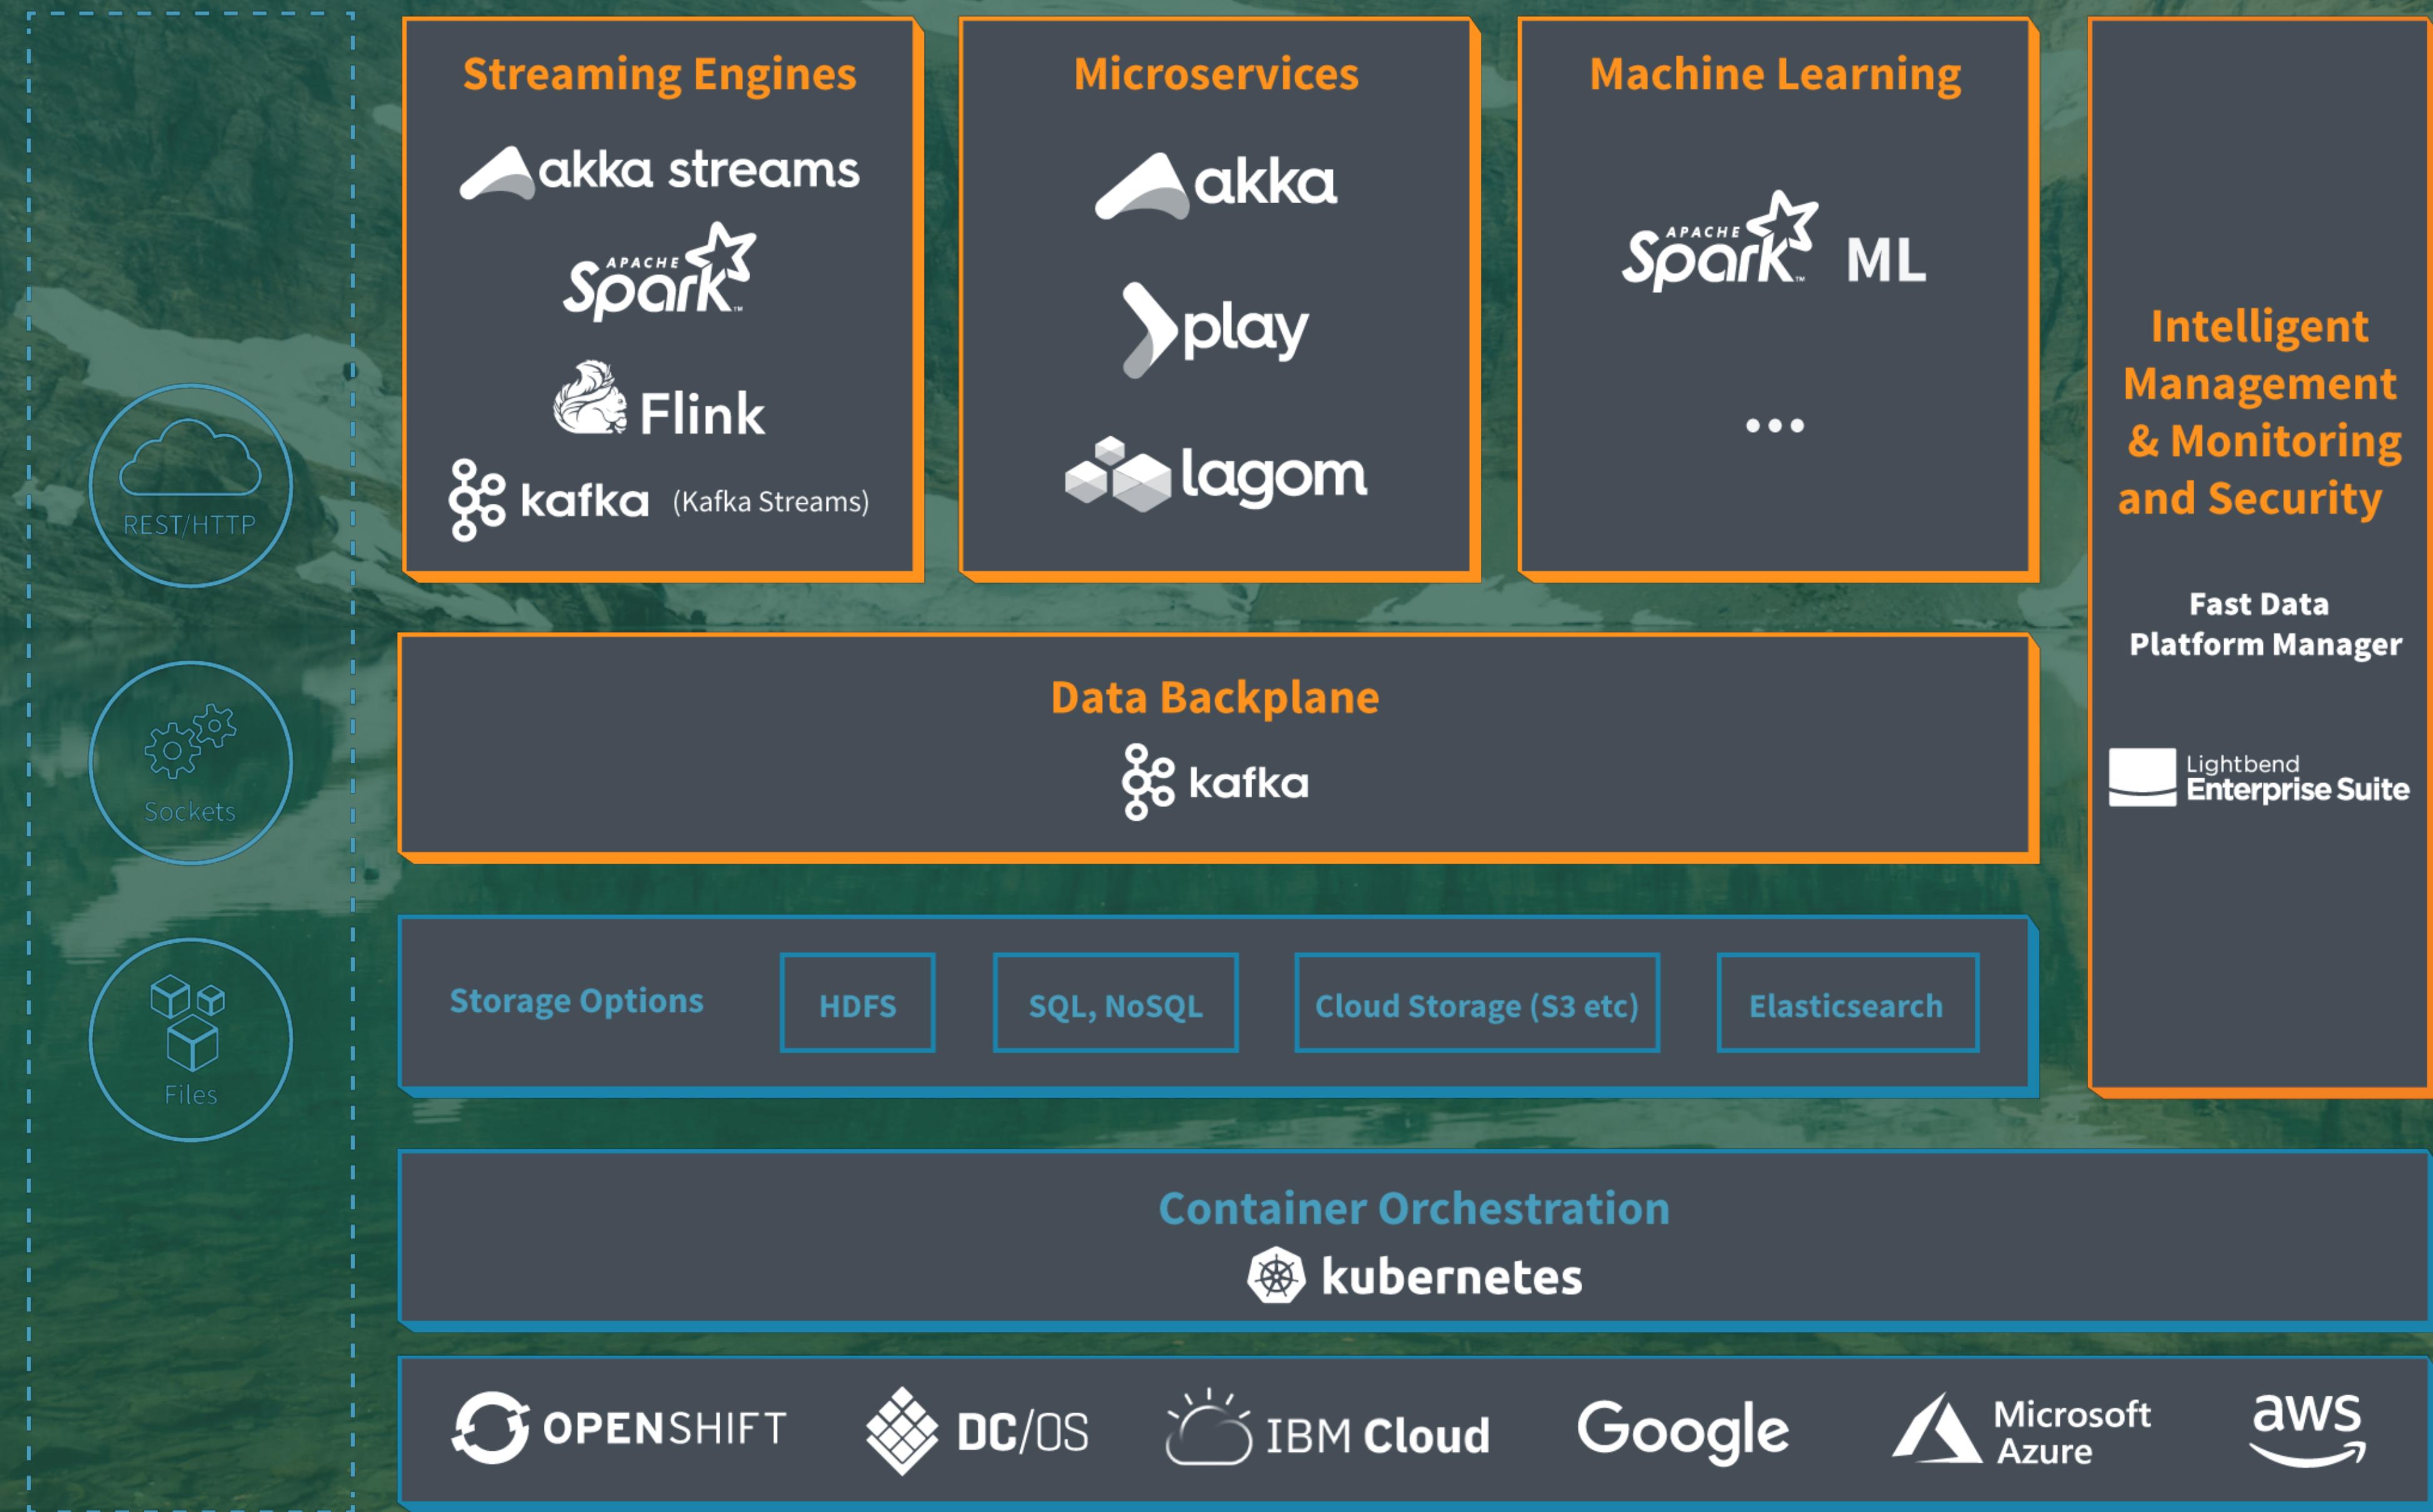
Getting Answers Now from  
Data Sets That Never End



2nd  
Edition

Dean Wampler, PhD

# lightbend.com/fast-data-platform





# Questions?

@deanwampler  
[lightbend.com/fast-data-platform](http://lightbend.com/fast-data-platform)  
[polyglotprogramming.com/talks](http://polyglotprogramming.com/talks)