

Why Big Data Needs to be Functional



Friday, March 9, 12

All pictures © Dean Wampler, 2011-2012.

What is Big Data?

Data so big that traditional solutions are too slow, too small, or too expensive to use.



It's a buzz word, but generally associated with the problem of data sets too big to manage with traditional SQL databases. A parallel development has been the NoSQL movement that is good at handling semistructured data, scaling, etc.



3 Trends

3

Friday, March 9, 12

Three trends influence my thinking...

Data Size ↑

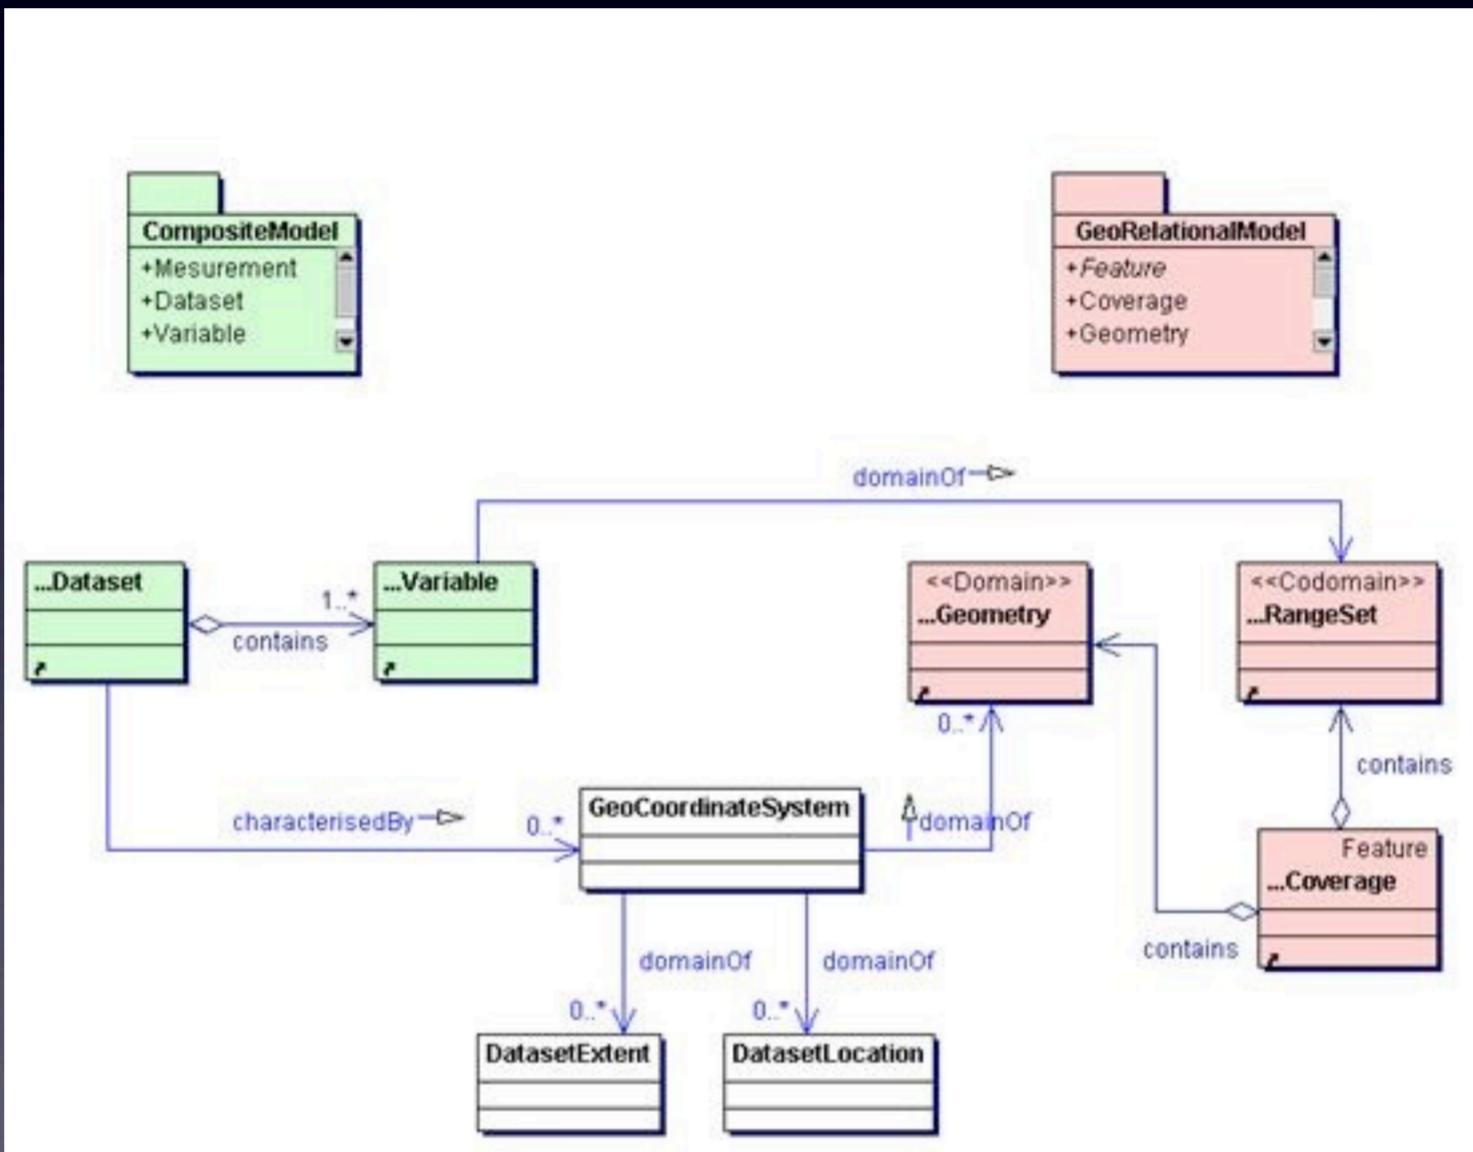


4

Friday, March 9, 12

Data volumes are obviously growing... rapidly.

Formal Schemas ↓



5

Friday, March 9, 12

There is less emphasis on “formal” schemas and domain models, because data changes rapidly, there are disparate data sources being joined, using relatively-agnostic software (e.g., collections of things where the software is agnostic about the contents) tends to be faster to develop and run.

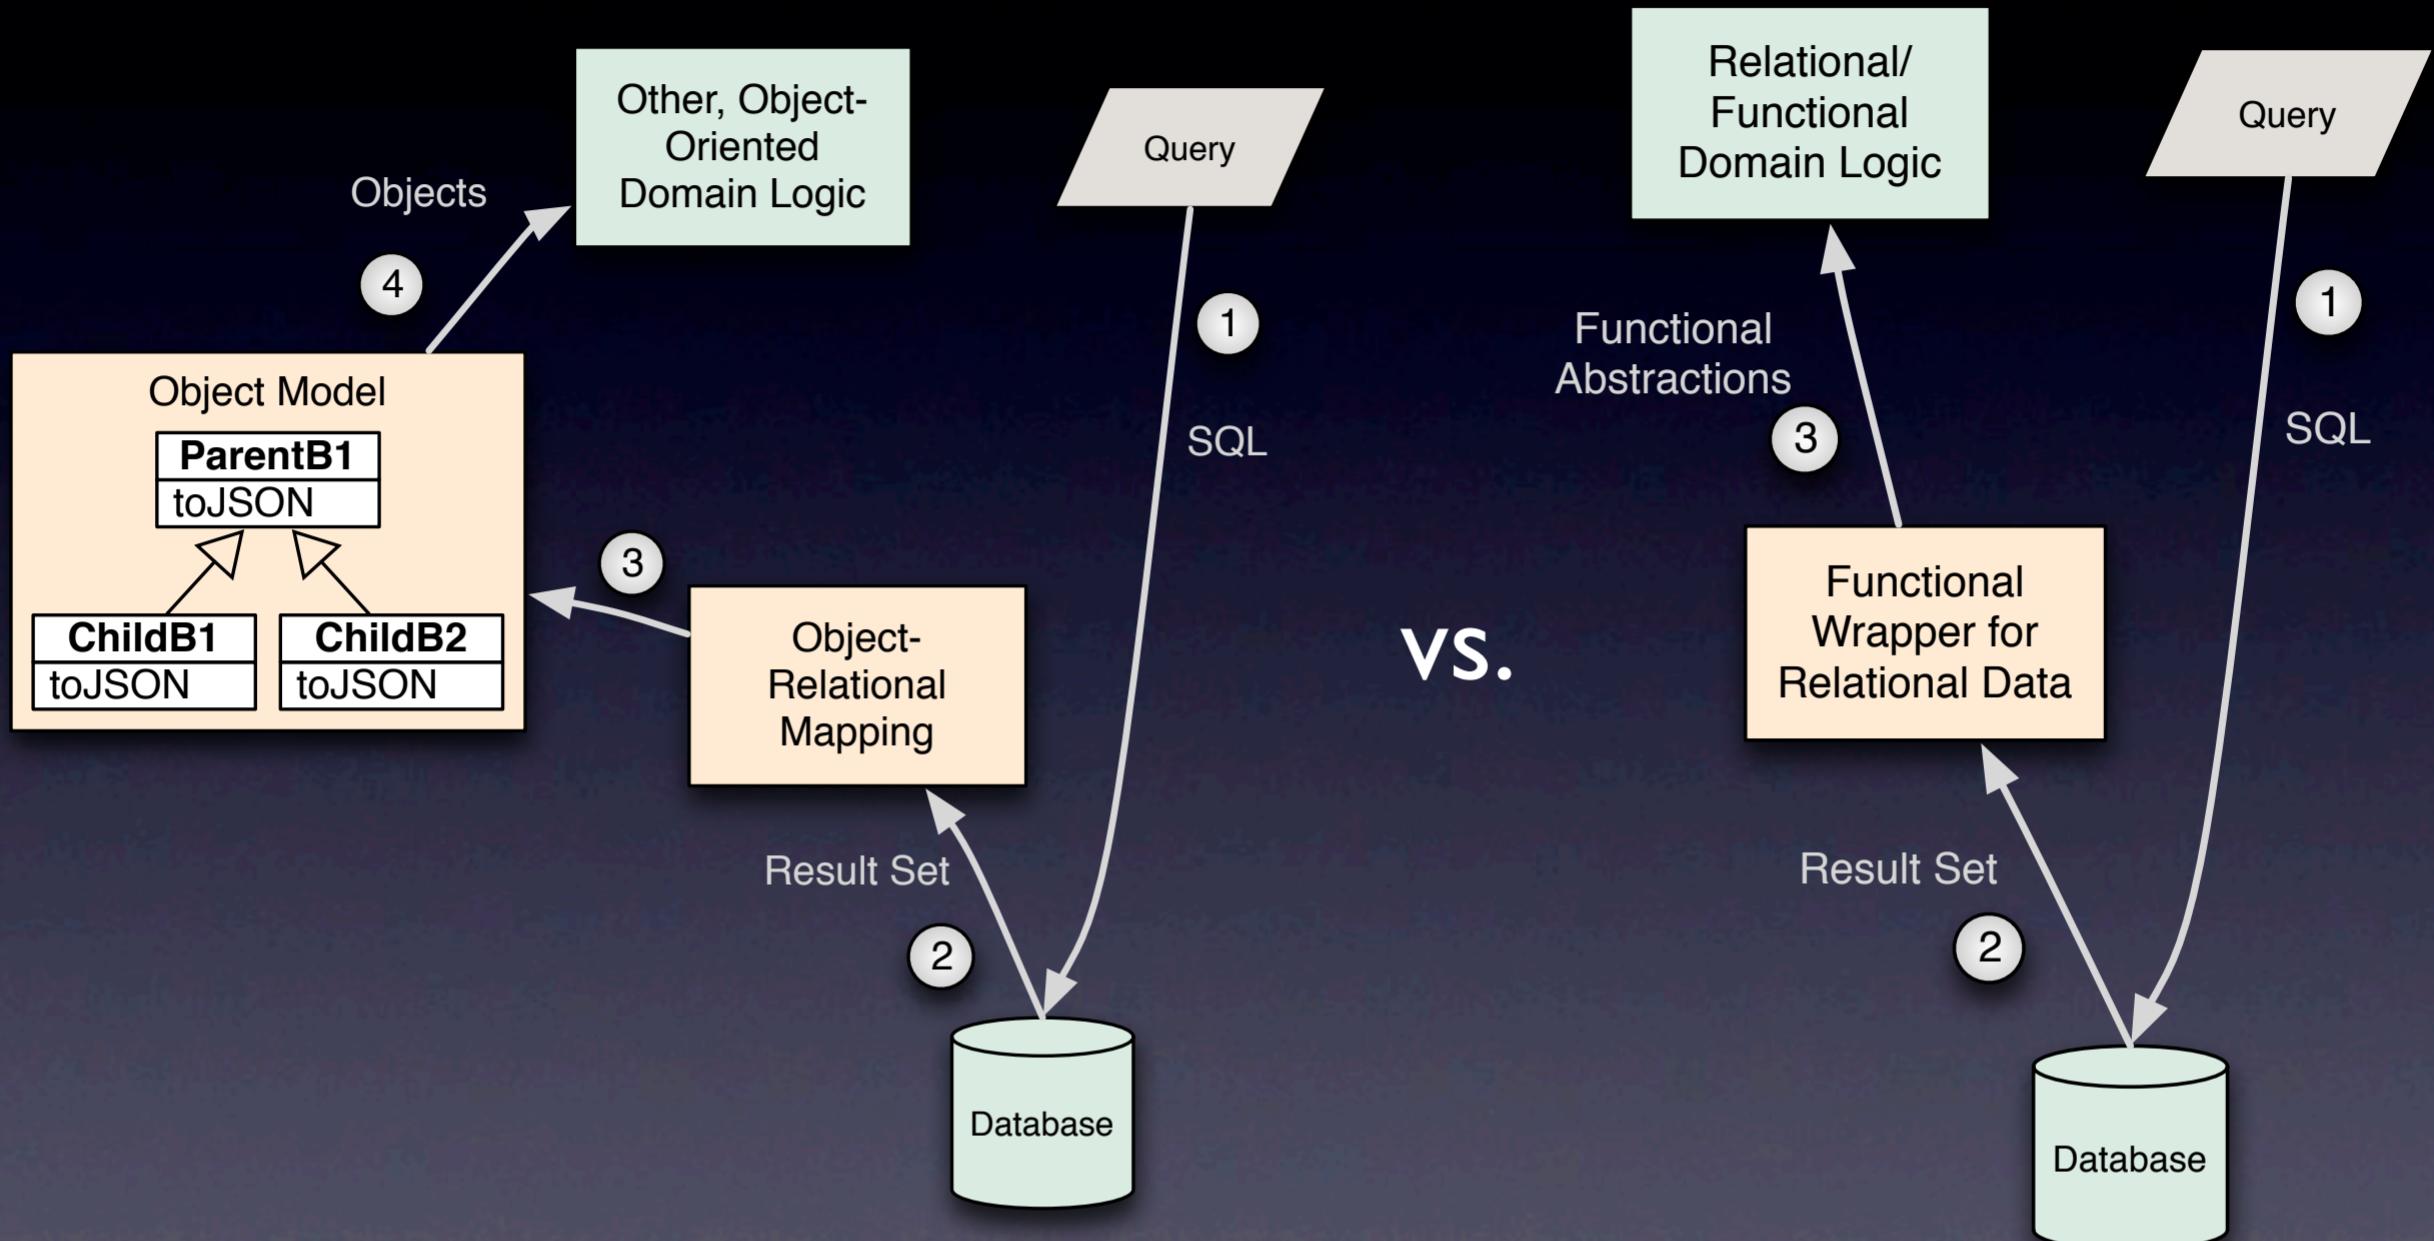
Data-Driven Programs ↑



6

Friday, March 9, 12

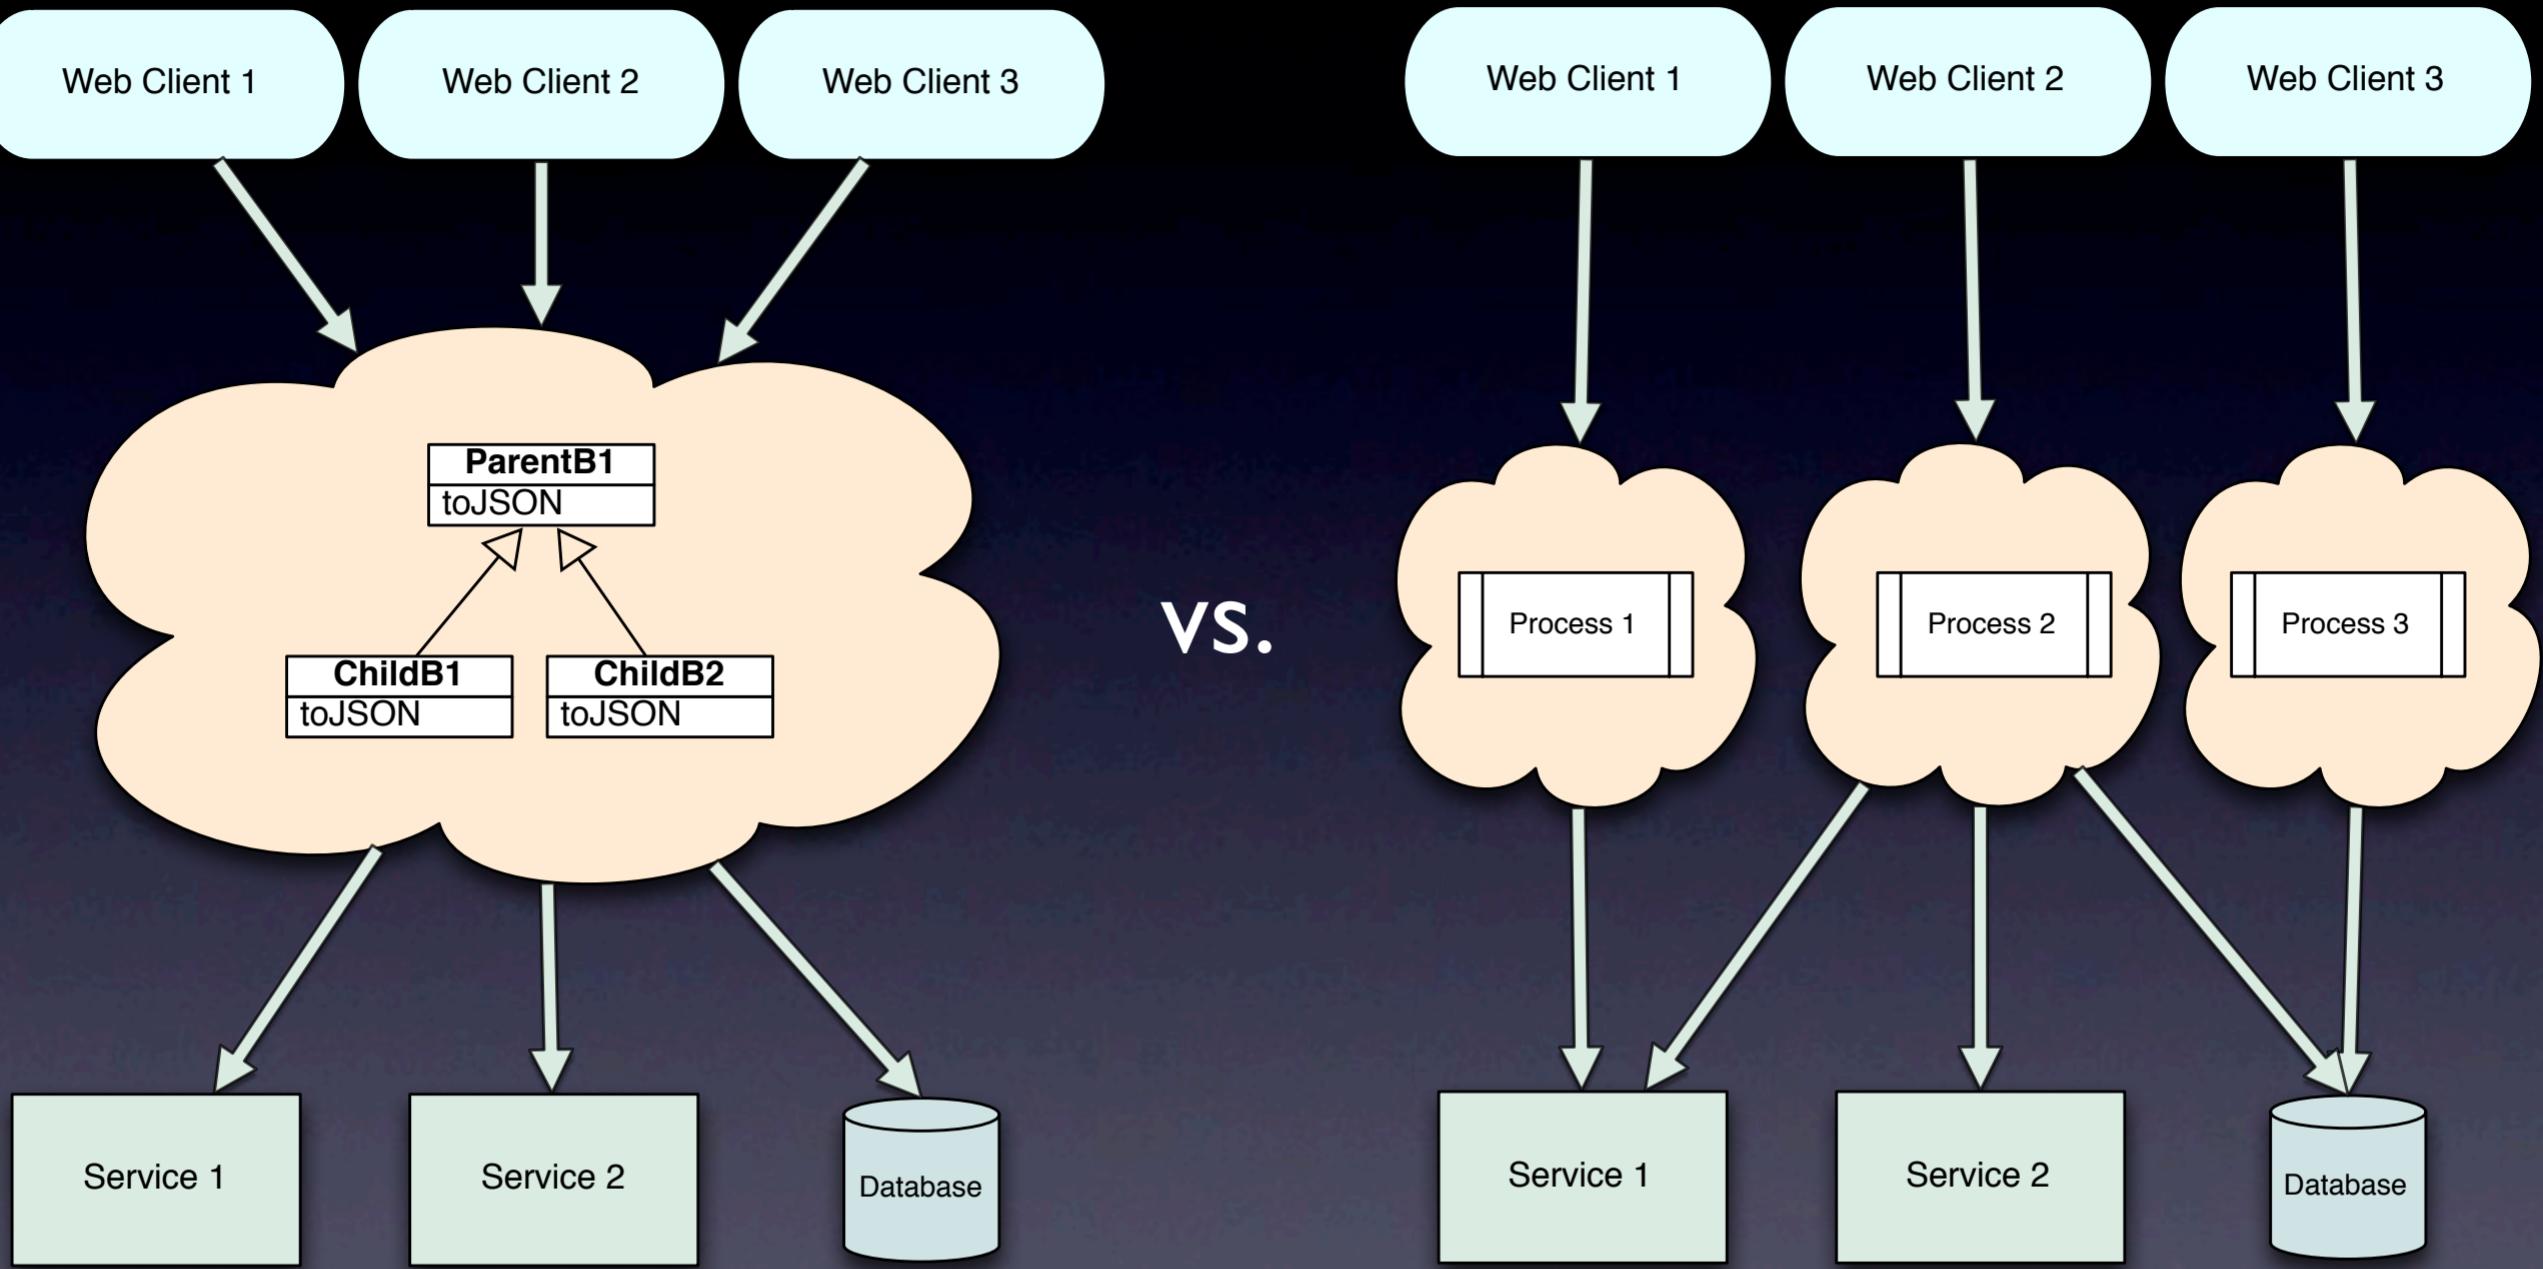
Machine learning is growing in importance. Here, generic algorithms and data structures are trained to represent the “world” using data, rather than encoding a model of the world in the software itself.



7

Friday, March 9, 12

Traditionally, (on the left) we've kept a rich, in-memory domain model requiring an ORM to convert persistent data into the model. This is resource overhead and complexity we can't afford in big data systems. Rather, (on the right) we should treat the result set as it is, a particular kind of collection, do the minimal transformation required to exploit our collections libraries and classes representing some domain concepts (e.g., Address, StockOption, etc.), then write functional code to implement business logic (or drive emergent behavior with machine learning algorithms...)

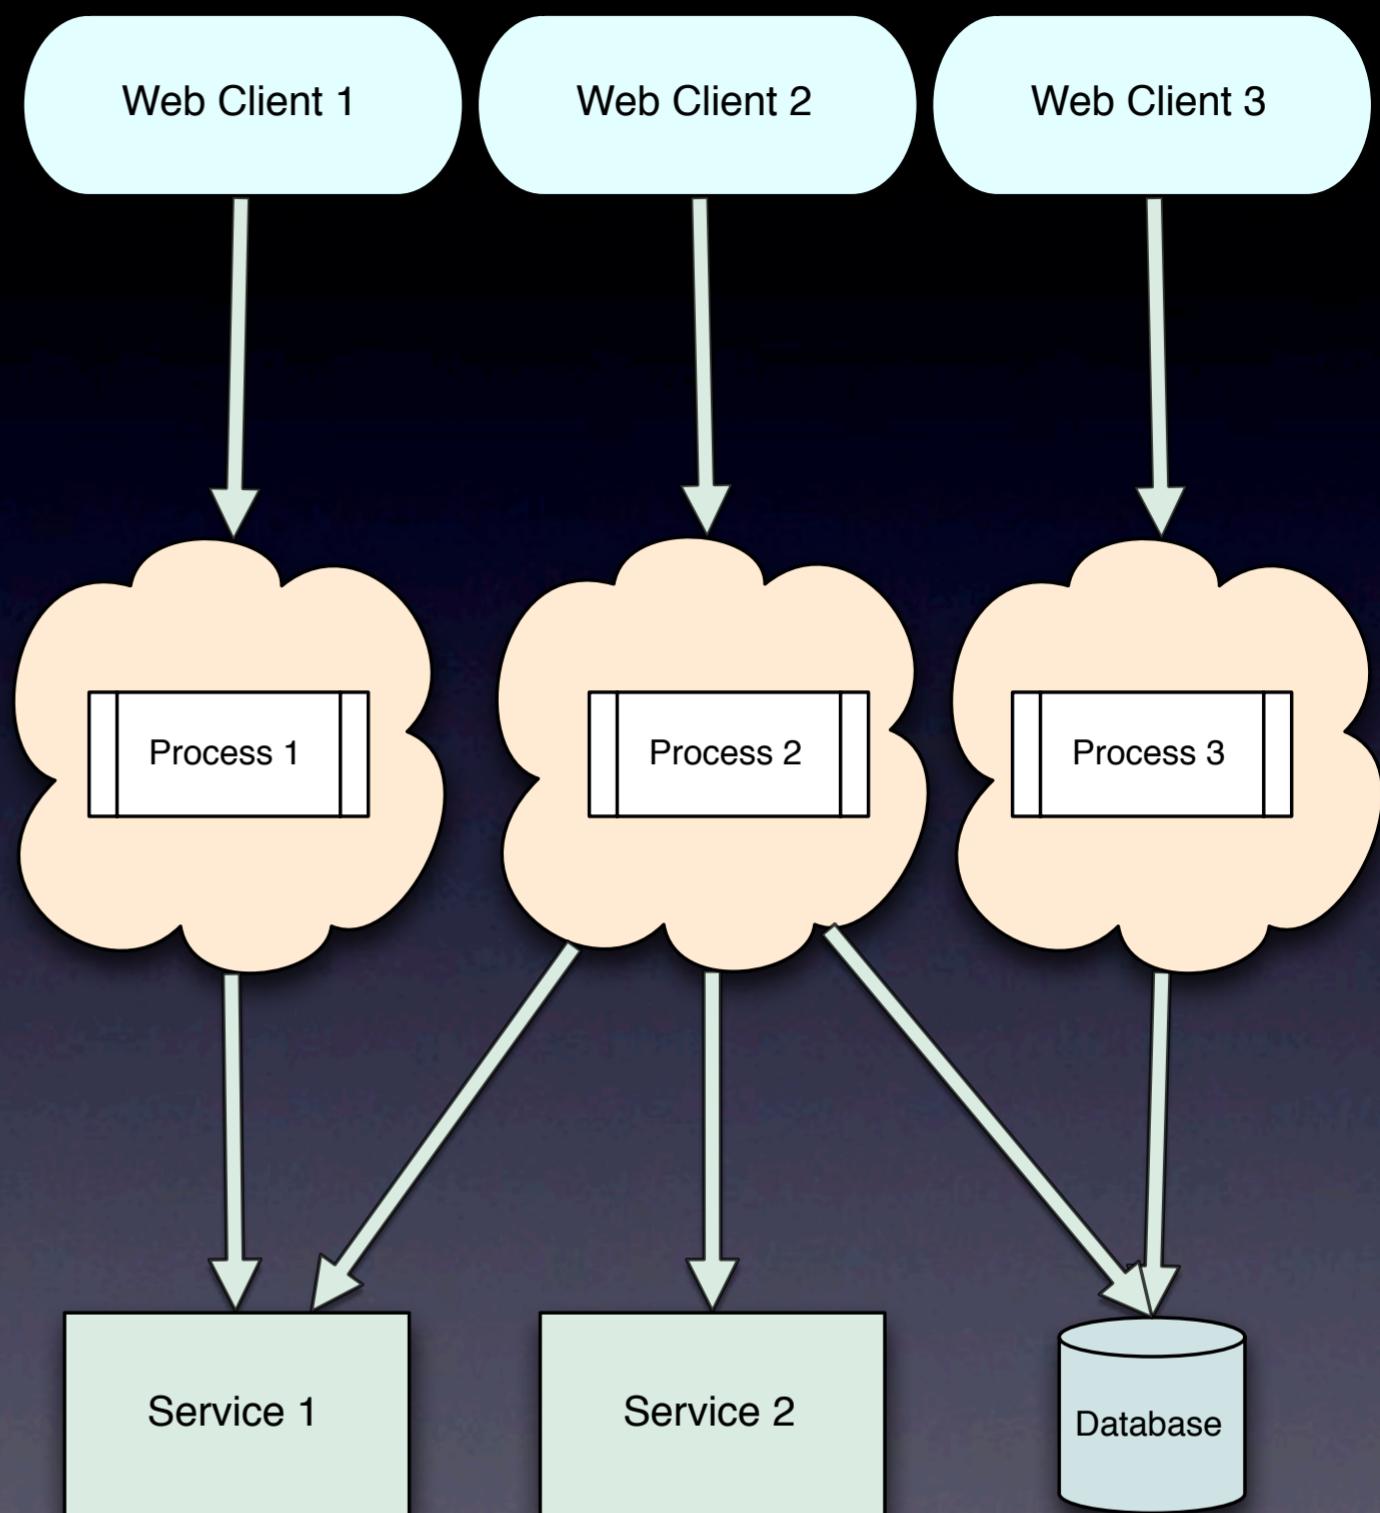


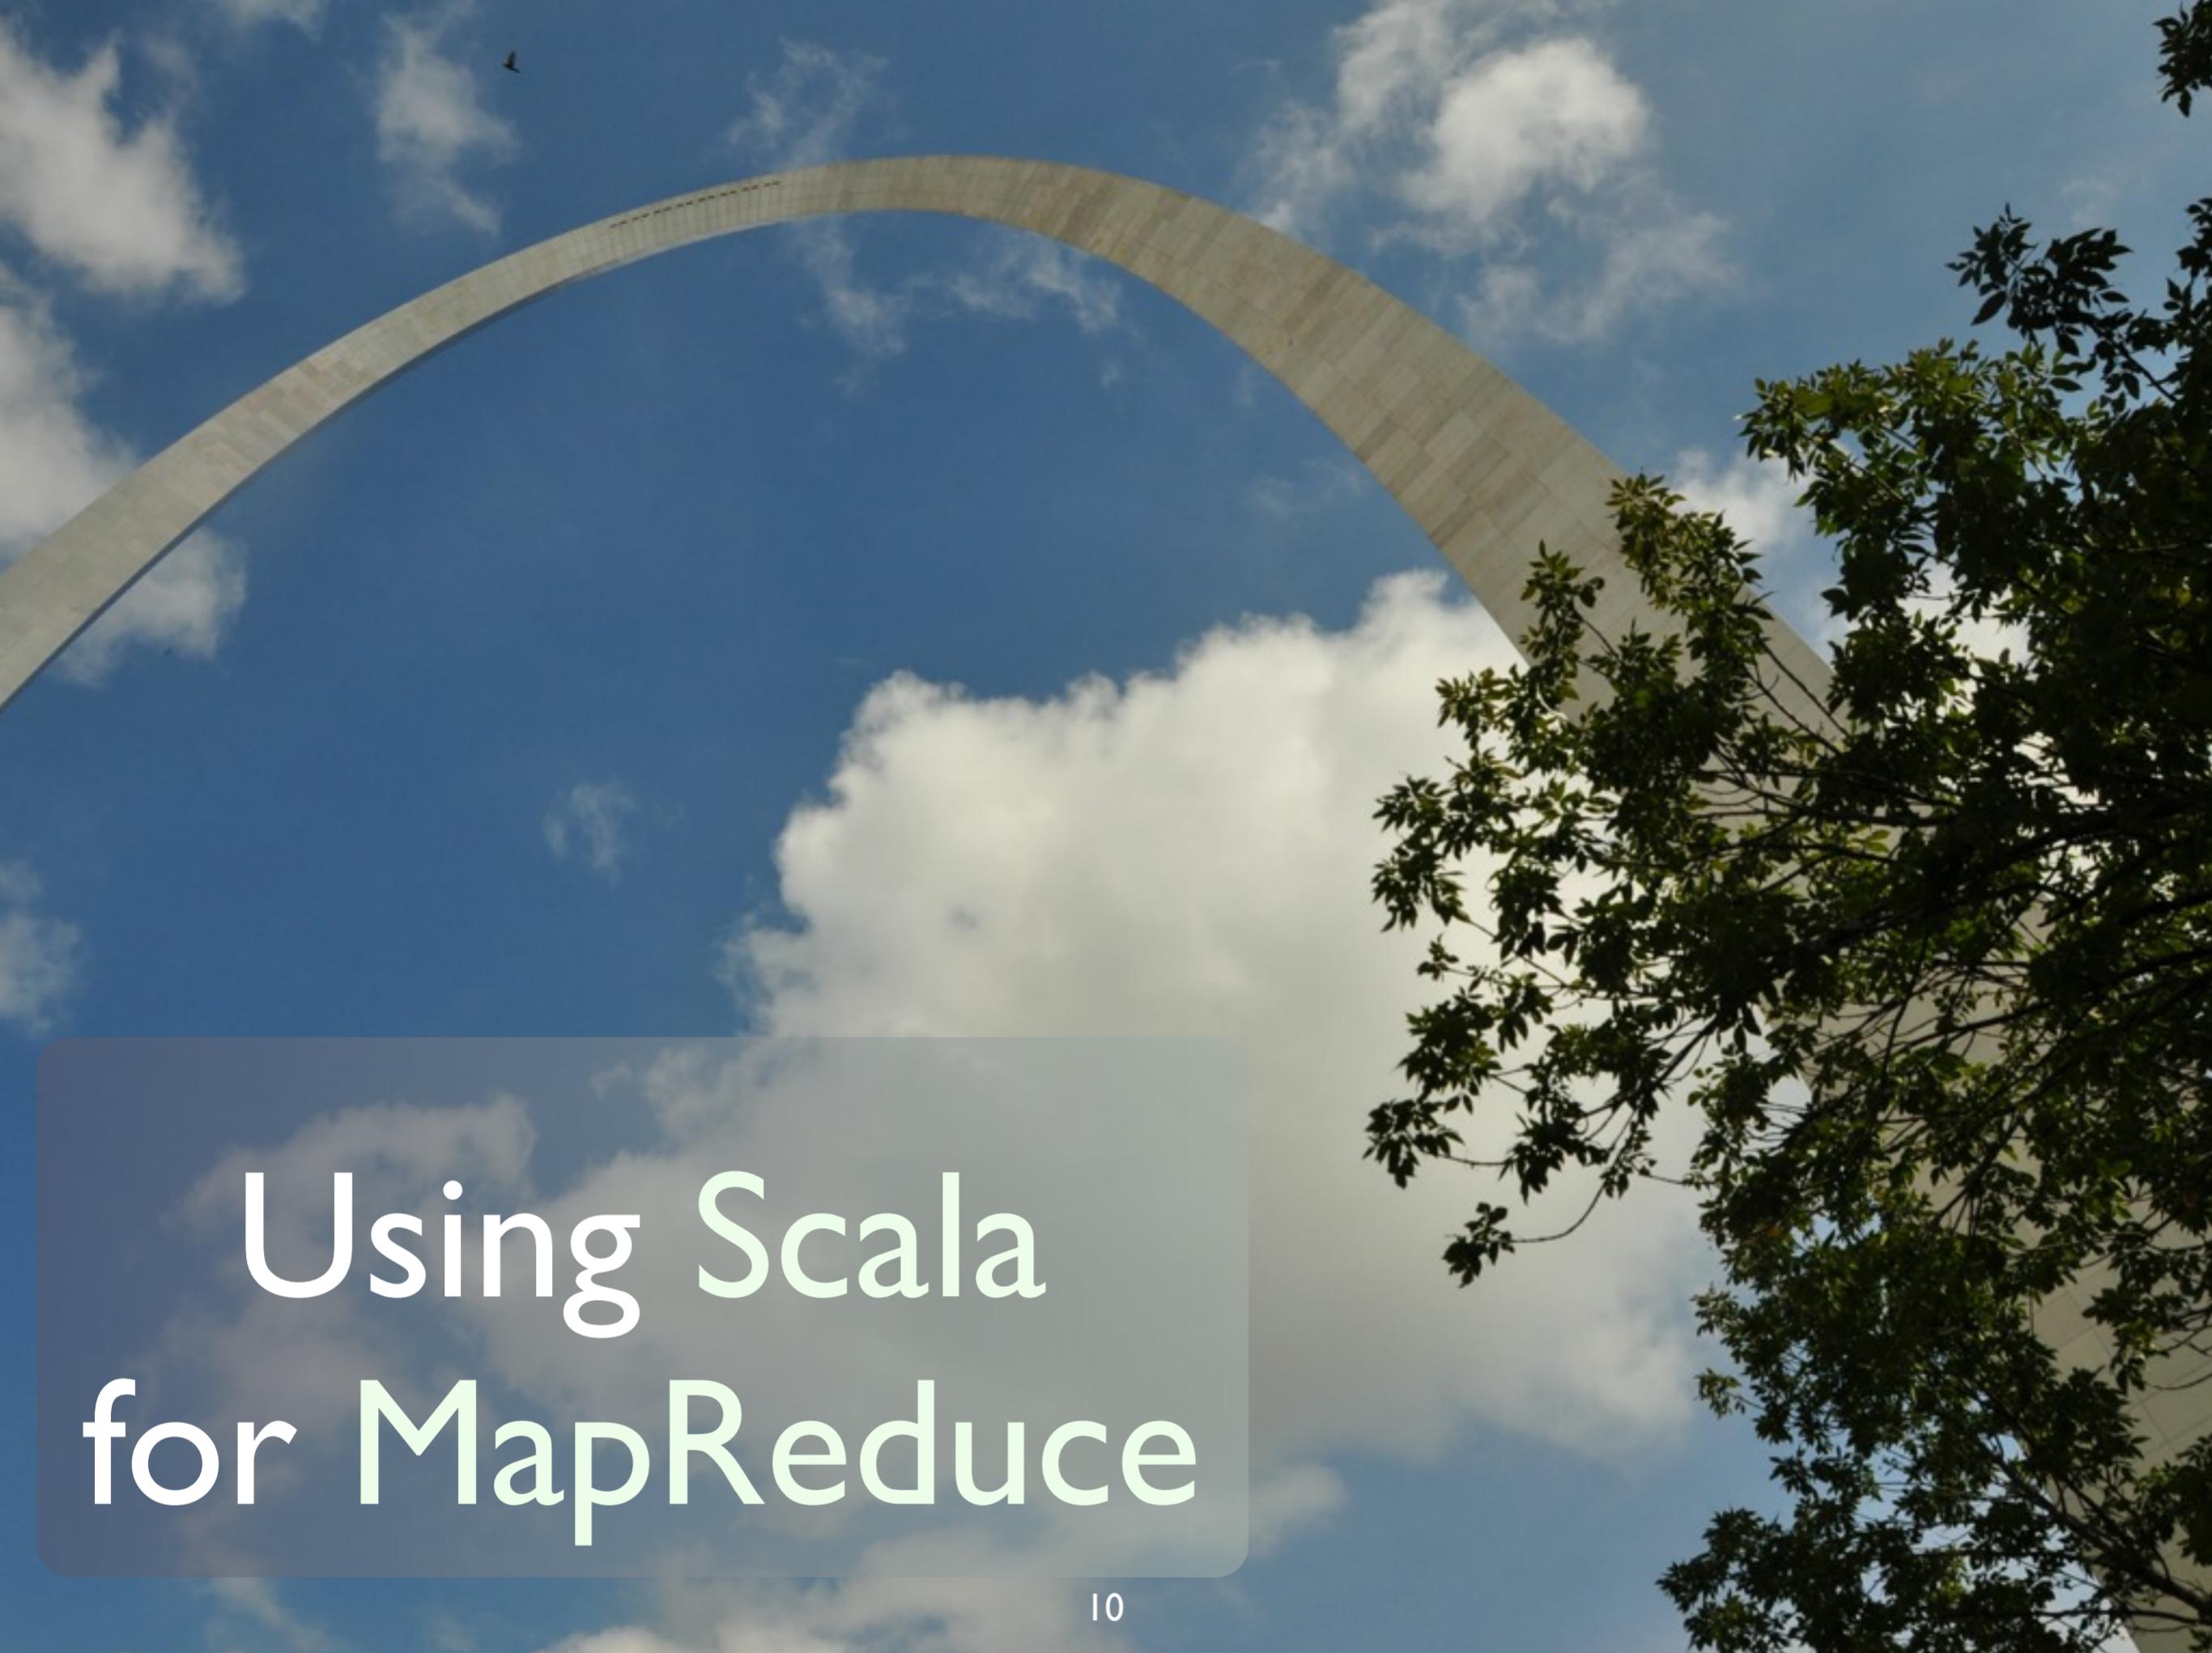
8

Friday, March 9, 12

In a broader view, object models (on the left) tend to push us towards centralized, complex systems that don't decompose well and stifle reuse and optimal deployment scenarios. FP code makes it easier to write smaller, focused services (on the right) that we compose and deploy as appropriate.

- Data Size ↑
- Formal Schema ↓
- Data-Driven Programs ↑





Using Scala for MapReduce

10

Friday, March 9, 12

Back to Scala, let's look at the mess that Java has introduced into Hadoop MapReduce, and get a glimpse of a better way. As an example, I'll walk you through the "Hello World" of MapReduce; the Word Count algorithm...

Input

Mappers



Consider Word Count

Friday, March 9, 12

Each document gets a mapper. I'm showing the document contents in the boxes for this example. Actually, large documents might get split to several mappers (as we'll see). It is also possible to concatenate many small documents into a single, larger document for input to a mapper.

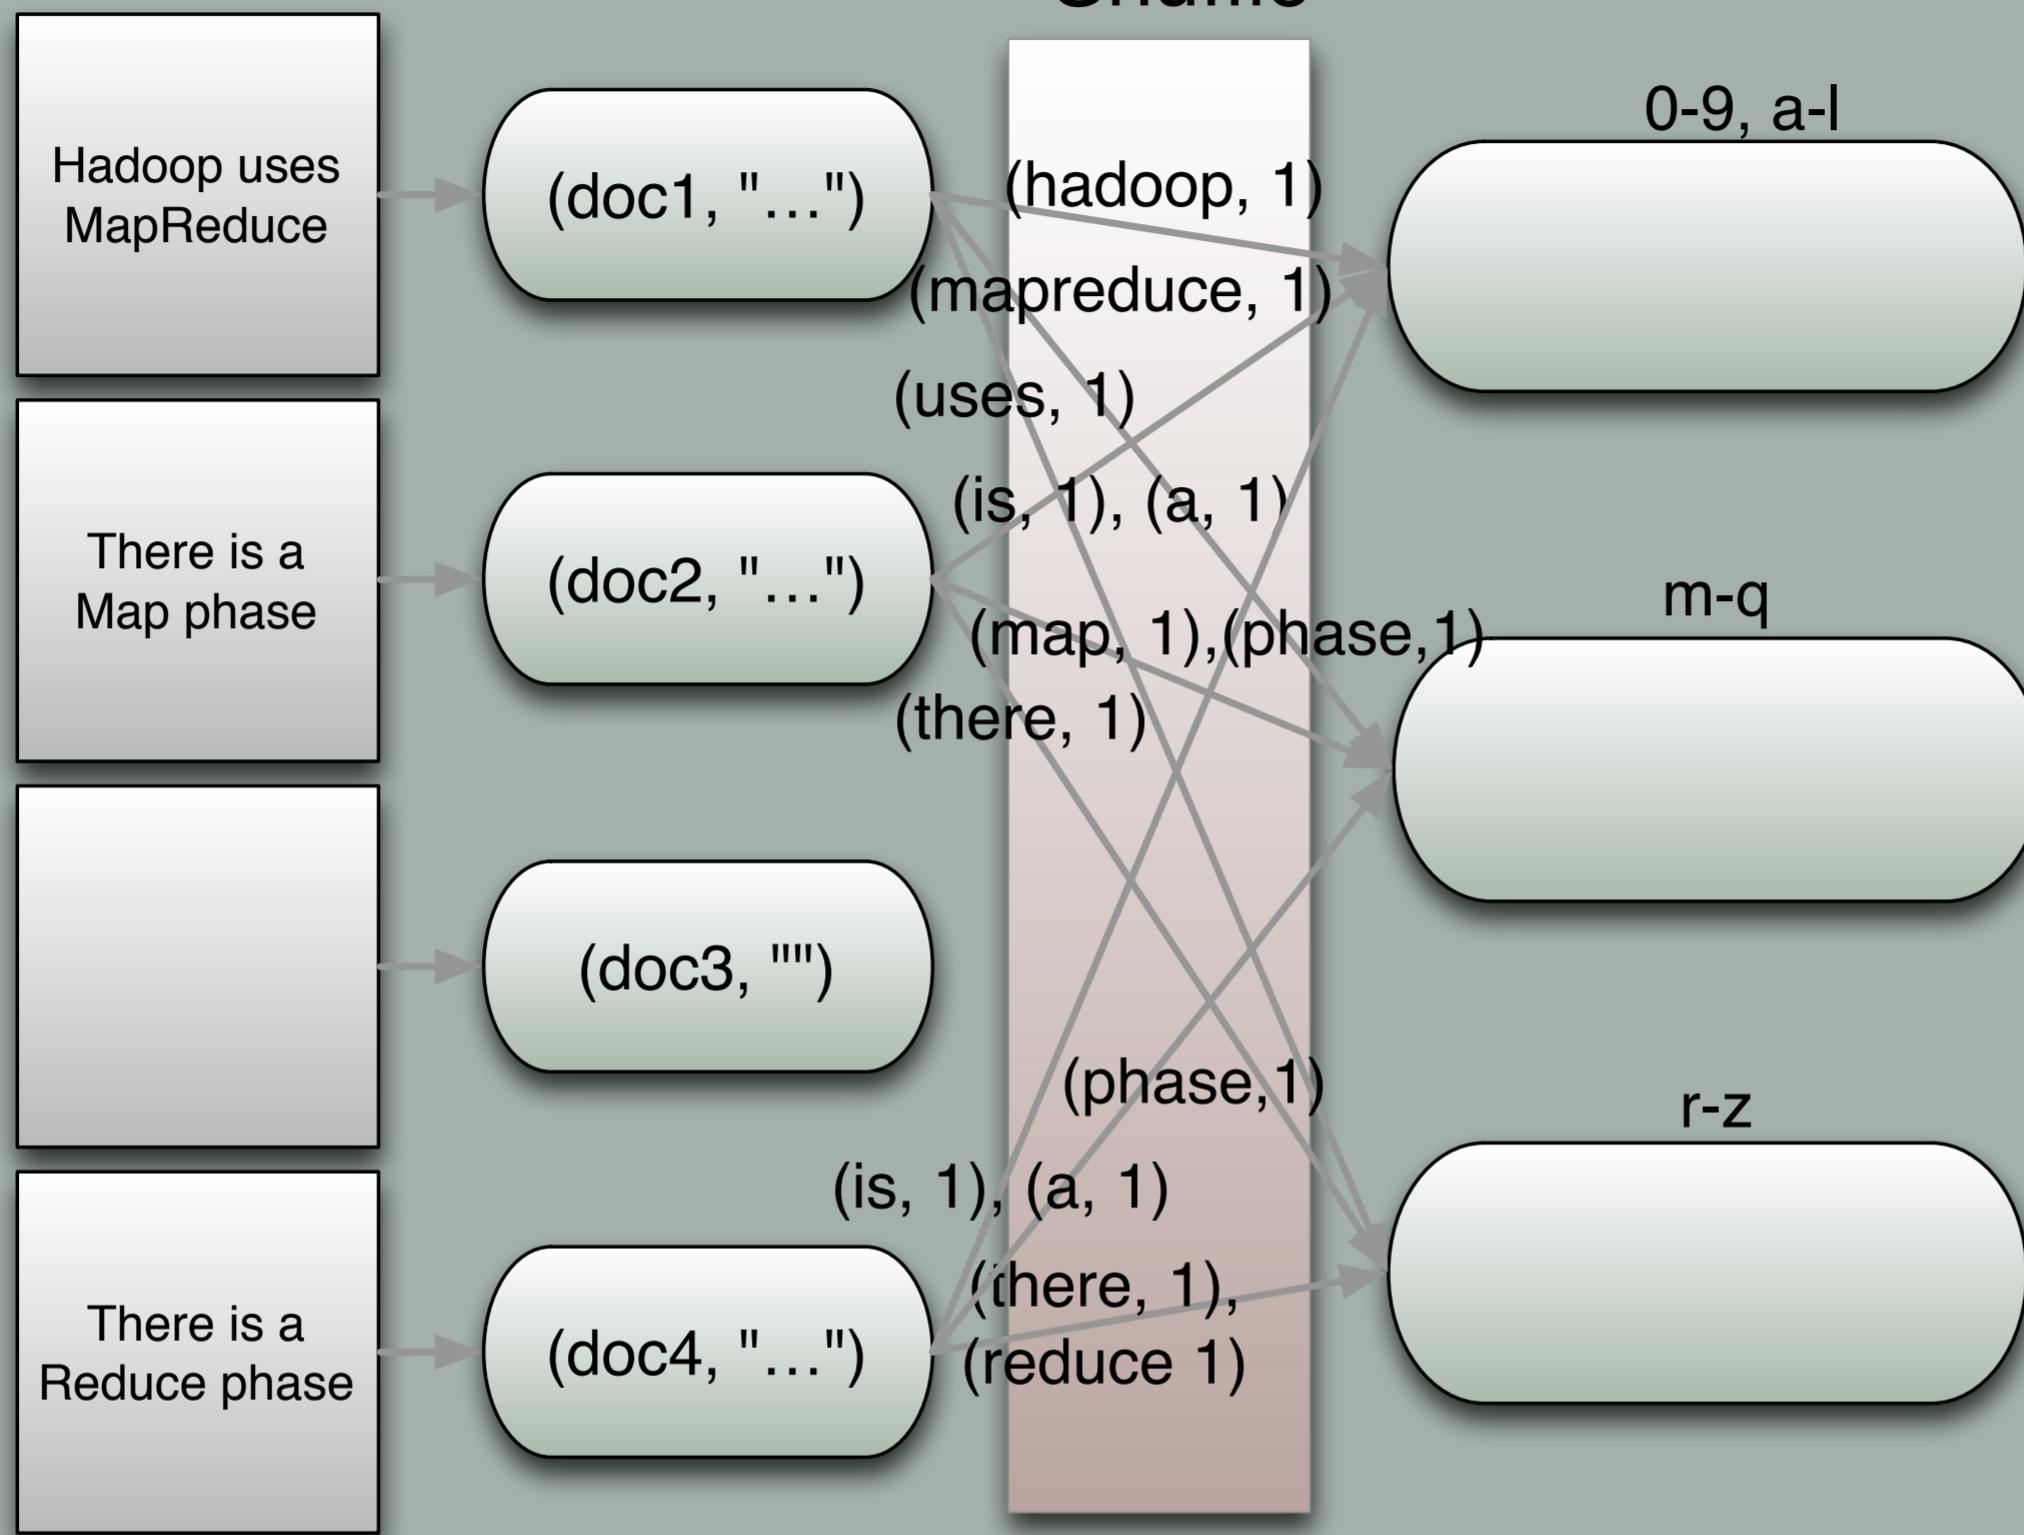
Each mapper will receive a key-value pair, where the key is the document path and the value is the contents of the document. It will ignore the key, tokenize the content, convert all words to lower case and count them...

Input

Mappers

Sort, Shuffle

Reducers



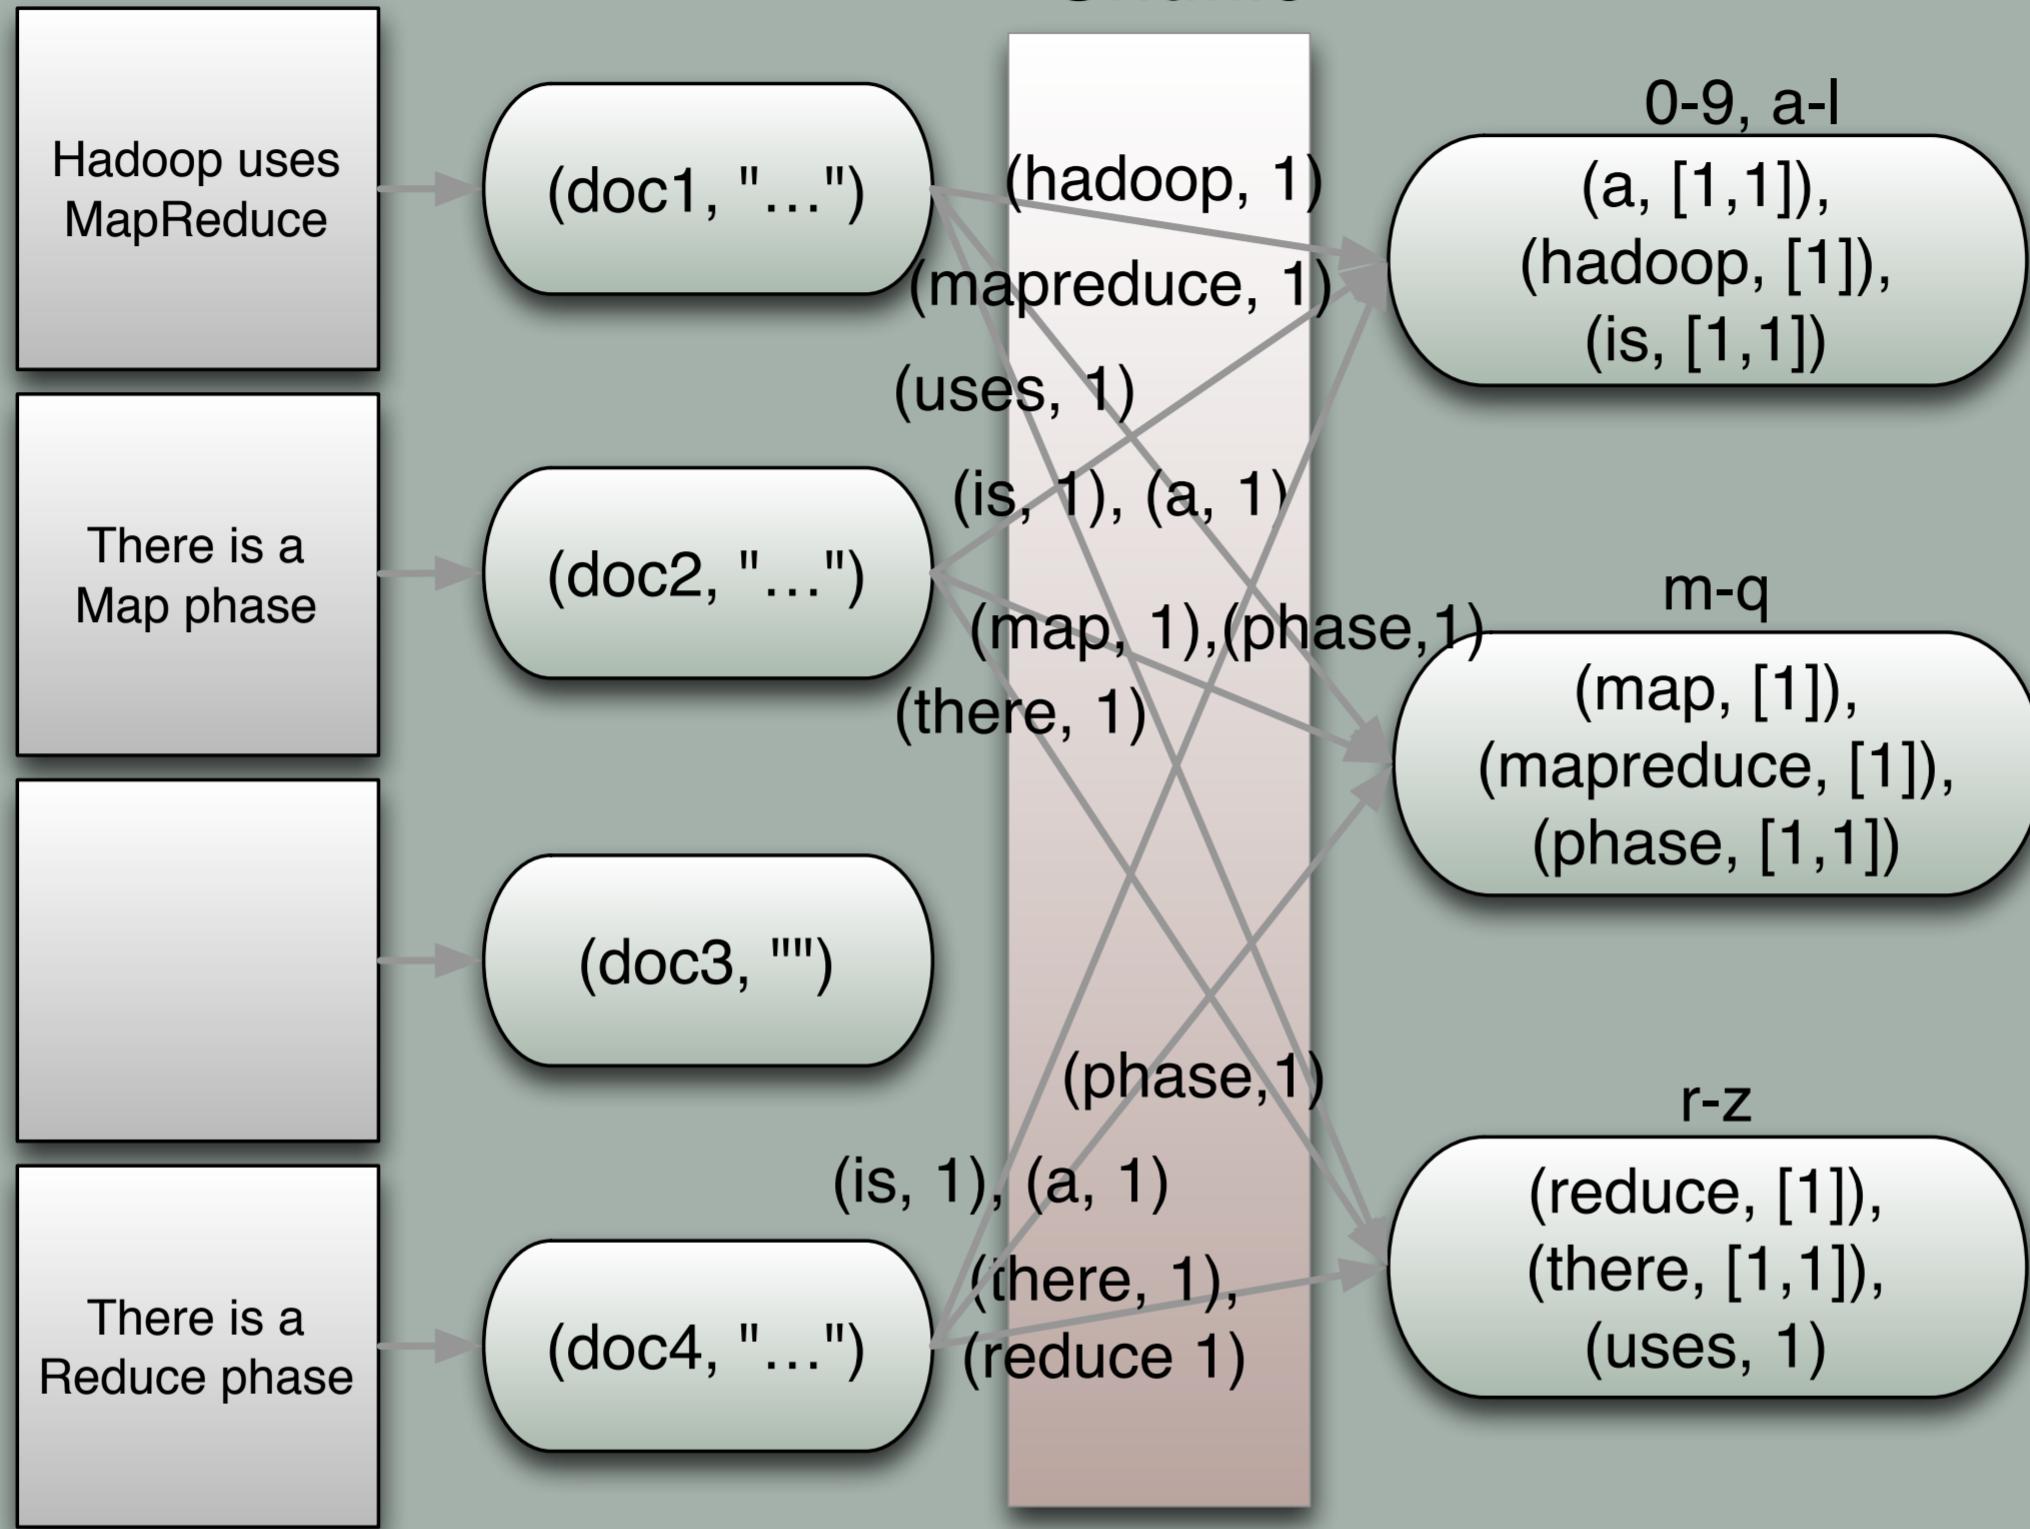
Friday, March 9, 12

The mappers emit key-value pairs, where each key is one of the words, and the value is the count. In the most naive (but also most memory efficient) implementation, each mapper simply emits (word, 1) each time “word” is seen.

The mappers themselves don’t decide to which reducer each pair should be sent. Rather, the job setup configures what to do and the Hadoop runtime enforces it during the Sort/Shuffle phase, where the key-value pairs in each mapper are sorted by key (that is locally, not globally or “totally”) and then the pairs are routed to the correct reducer, on the current machine or other machines.

Note how we partitioned the reducers (by first letter of the keys). Also, note that the mapper for the empty doc. emits no pairs, as you would expect.

Input Mappers Sort, Shuffle Reducers

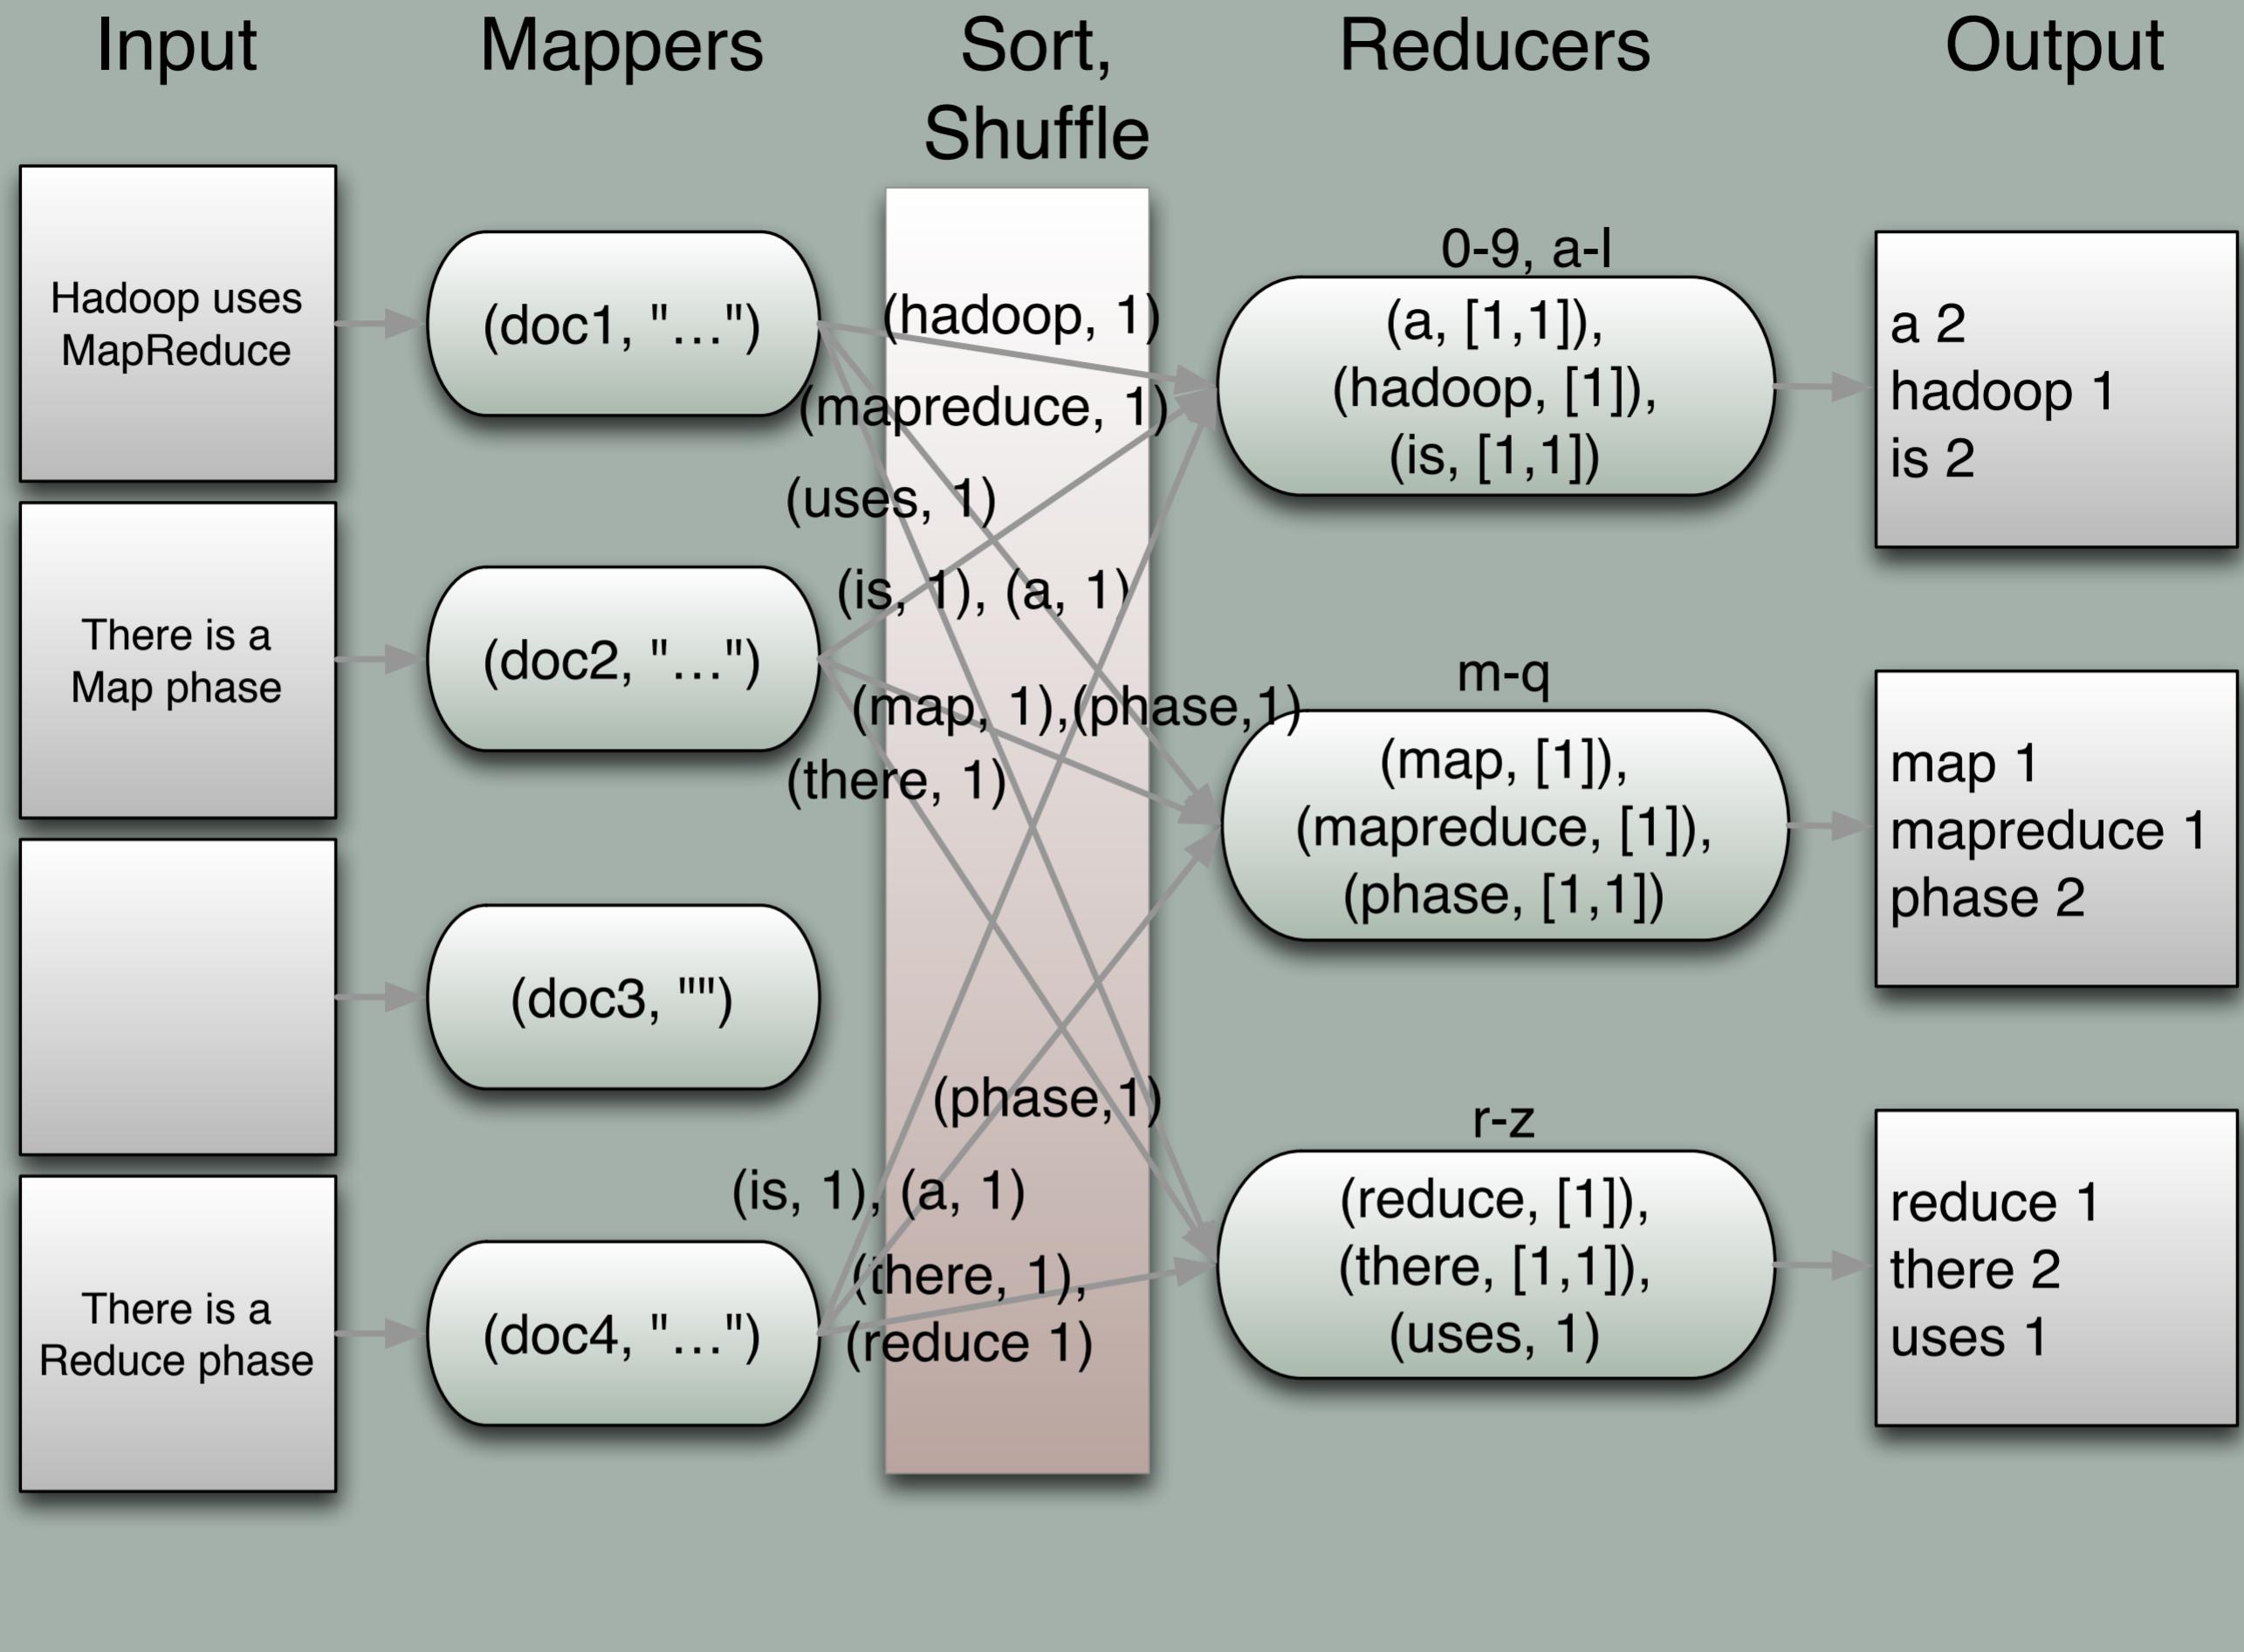


Friday, March 9, 12

The mappers emit key-value pairs, where each key is one of the words, and the value is the count. In the most naive (but also most memory efficient) implementation, each mapper simply emits (word, 1) each time “word” is seen.

The mappers themselves don’t decide to which reducer each pair should be sent. Rather, the job setup configures what to do and the Hadoop runtime enforces it during the Sort/Shuffle phase, where the key-value pairs in each mapper are sorted by key (that is locally, not globally or “totally”) and then the pairs are routed to the correct reducer, on the current machine or other machines.

Note how we partitioned the reducers (by first letter of the keys). Also, note that the mapper for the empty doc. emits no pairs, as you would expect.



Friday, March 9, 12

The final view of the WordCount process flow for our example. We'll see in more detail shortly how the key-value pairs are passed to the reducers, which add up the counts for each word (key) and then writes the results to the output files. The output files contain one line for each key (the word) and value (the count), assuming we're using text output. The choice of delimiter between key and value is up to you. (We'll discuss options as we go.)

The MapReduce Java API

15

Friday, March 9, 12

A Java API, but I'll show you Scala code; see my GitHub `scala-hadoop` project.

```

import org.apache.hadoop.io.*
import org.apache.hadoop.mapred.*
import java.util.StringTokenizer

object SimpleWordCount {

  val one = new IntWritable(1)
  val word = new Text      // Value will be set in a non-thread-safe way!

  class WCMapper extends MapReduceBase with Mapper[LongWritable, Text, Text, IntWritable] {

    def map(key: LongWritable, valueDocContents: Text,
           output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {
      val tokens = valueDocContents.toString.split("\\s+")
      for (wordString <- tokens) {
        if (wordString.length > 0) {
          word.set(wordString.toLowerCase)
          output.collect(word, one)
        }
      }
    }
  }

  class Reduce extends MapReduceBase with Reducer[Text, IntWritable, Text, IntWritable] {

    def reduce(keyWord: Text, valuesCounts: java.util.Iterator[IntWritable],
               output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {
      var totalCount = 0
      while (valuesCounts.hasNext) {
        totalCount += valuesCounts.next.get
      }
      output.collect(keyWord, new IntWritable(totalCount))
    }
  }
}

```

Using Crunch (Java)



17

Friday, March 9, 12

Crunch is a Java library that provides a higher-level abstraction for data computations and flows on top of MapReduce, inspired by a paper from Google researchers on an internal project called FlumeJava.

See <https://github.com/cloudera/crunch>.

Crunch Concepts

- *Pipeline<T>*: Abstraction for data processing flow.
- *PCollection<T>*: Distributed, unordered collection of elements T.
- *parallelDo()*: Apply operation across collection.

Crunch Concepts

- *PTable<K,V>*: Distributed multimap.
 - *groupByKey()*: group together all values with the same key.
 - *parallelDo()*: apply operation across collection.

Crunch Concepts

- *PGroupedTable<K,V>*: Output of *groupByKey()*.
 - *combineValues()*: associative and commutative .
 - *groupByKey()*: group together all values with the same key.
 - *parallelDo()*: apply operation across collection.

```

import com.cloudera.crunch.*;
import org.apache.hadoop.*;

...
public class WordCount extends Configured implements Tool, Serializable {
    public int run(String[] args) throws Exception {
        Pipeline pipeline = new MRPipeline(WordCount.class, getConf());
        PCollection<String> lines = pipeline.readTextFile(args[0]);

        PCollection<String> words = lines.parallelDo(new DoFn<String, String>() {
            public void process(String line, Emitter<String> emitter) {
                for (String word : line.split("\\s+")) {
                    emitter.emit(word);
                }
            }
        }, Writables.strings());
        PTable<String, Long> counts = Aggregate.count(words);
        pipeline.writeTextFile(counts, args[1]);
        pipeline.done();
        return 0;
    }
}

```

Using Scrunch (Scala)



22

Friday, March 9, 12

Scrunch is a scala DSL around Crunch writing by the same developers at Cloudera and also included in the Crunch distro. It uses type classes and similar constructs to provide wrap the Crunch classes and MR classes with REAL map, flatMap, etc. functionality.

```

import com.cloudera.crunch.*;
import com.cloudera.scrunch.*;
...
class ScrunchWordCount {
    def wordCount(inputFile: String,
                  outputFile: String) = {
        val pipeline = new Pipeline[WordCountExample]
        pipeline.read(from.textFile(inputFile))
            .flatMap(_.toLowerCase.split("\\w+"))
            .filter(!_isEmpty())
            .count
            .write(to.textFile(outputFile)) // Word counts
            .map((w, c) => (w.slice(0, 1), c))
            .groupByKey.combine(v => v.sum).materialize
        pipeline.done
    }
}
object ScrunchWordCount {
    def main(args: Array[String]) = {
        new ScrunchWordCount.wordCount(args(0), args(1))
    }
}

```

23

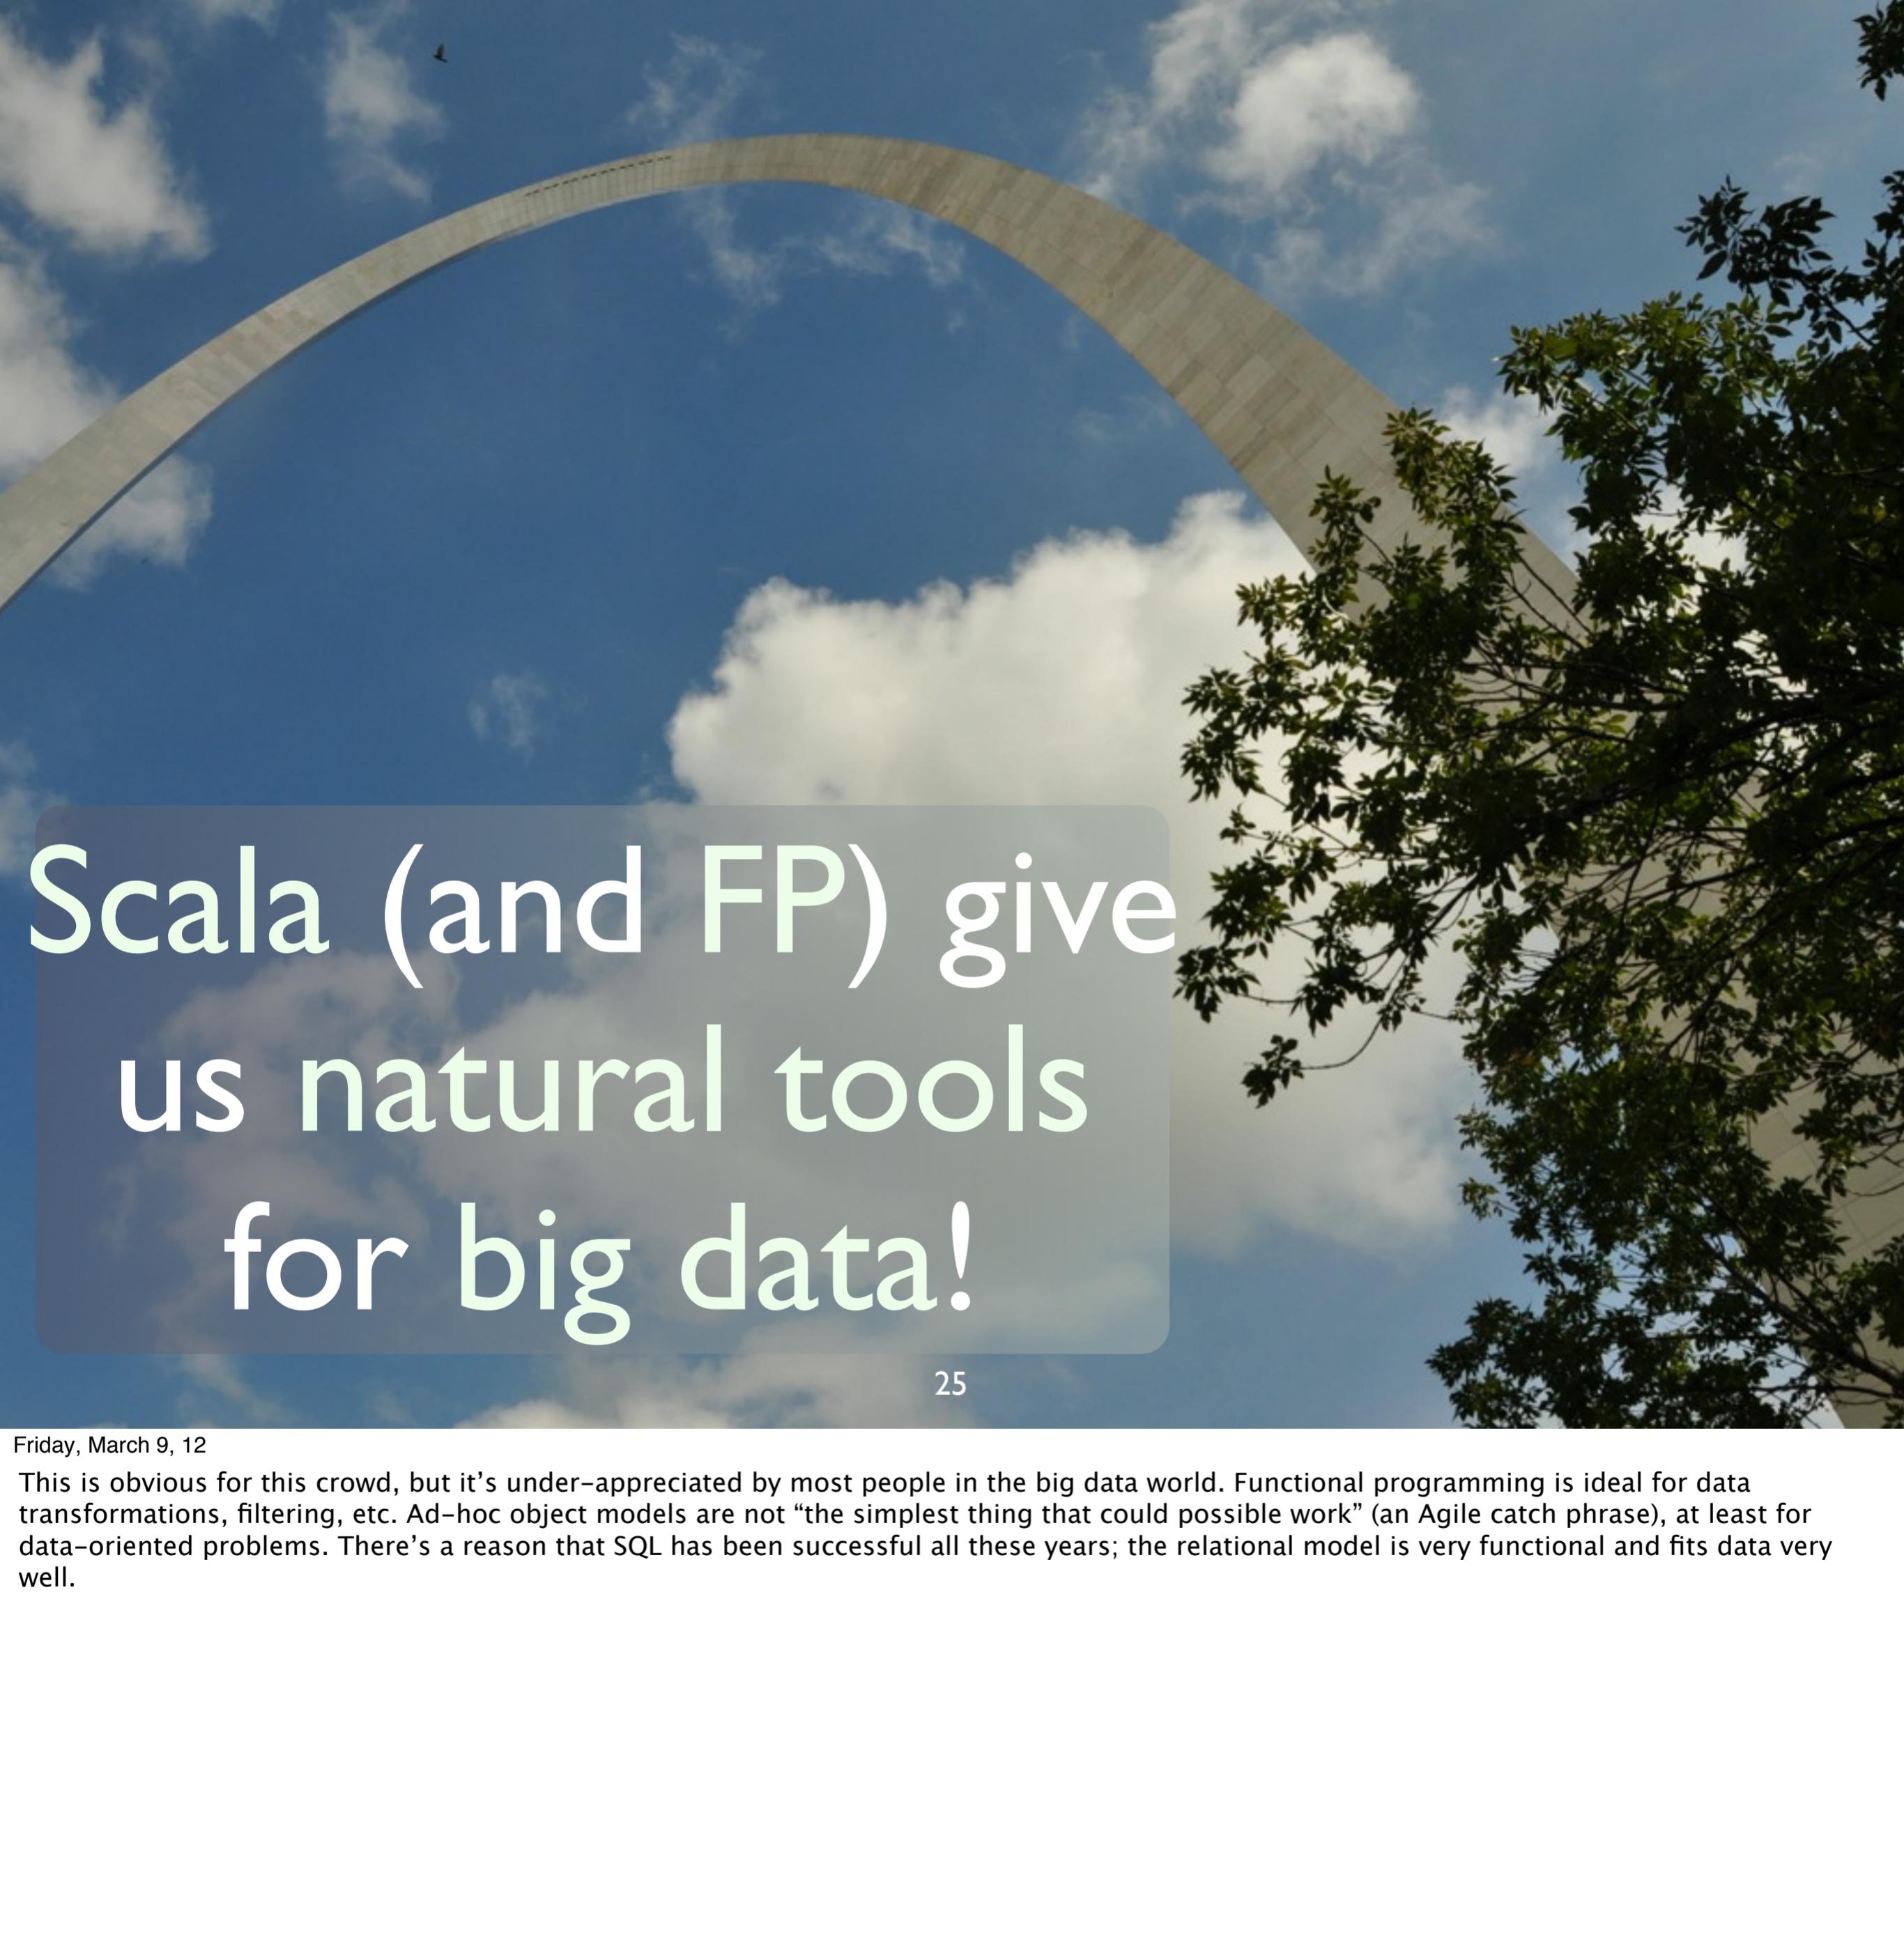
Friday, March 9, 12

(Back to Scala) I cheated; I'm showing you the WHOLE program, "main" and all. Not only is the size smaller and more concise still, but note the "builder" style notation, which intuitive lays out the data flow required. Note there is less green - fewer types are explicitly tossed about, and not just because Scala does some implicit typing. In contrast, there is more yellow - function calls showing the sequence of operations that more naturally represent the real "business logic", that is Word Count!

Also, fewer comments are required to help you sort out what's going on.

You Also See This with...

- *Cascading* (Java) vs.
 - *Cascalog* (Clojure)
 - *Scalding* (Scala)



Scala (and FP) give us natural tools for big data!

25

Friday, March 9, 12

This is obvious for this crowd, but it's under-appreciated by most people in the big data world. Functional programming is ideal for data transformations, filtering, etc. Ad-hoc object models are not "the simplest thing that could possibly work" (an Agile catch phrase), at least for data-oriented problems. There's a reason that SQL has been successful all these years; the relational model is very functional and fits data very well.



A Manifesto...

26

Friday, March 9, 12

So, I think we have an opportunity...



Hadoop is the Enterprise Java Beans of our time.

Friday, March 9, 12

I worked with EJBs a decade ago. The framework was completely invasive into your business logic. There were too many configuration options in XML files. The framework “paradigm” was a poor fit for most problems (like soft real time systems and most algorithms beyond Word Count). Internally, EJB implementations were inefficient and hard to optimize, because they relied on poorly considered object boundaries that muddled more natural boundaries. (I’ve argued in other presentations and my “FP for Java Devs” book that OOP is a poor modularity tool...) The fact is, Hadoop reminds me of EJBs in almost every way. It works okay and people do get stuff done, but just as the Spring Framework brought an essential rethinking to Enterprise Java, I think there is an essential rethink that needs to happen in Big Data. The Scala community is well positioned to create it.



Scala Collections.

Friday, March 9, 12

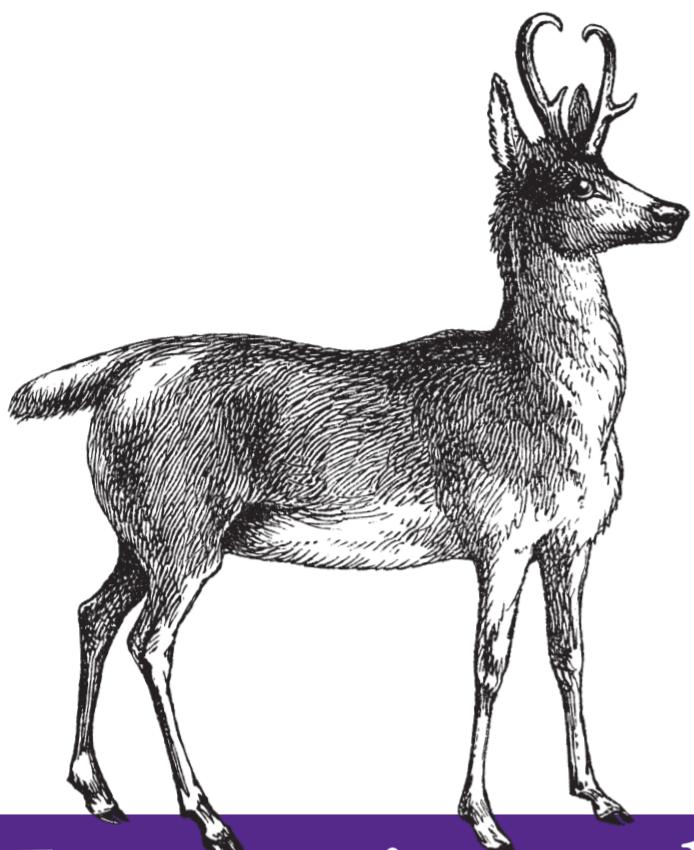
We already have the write model in Scala's collections and the parallel versions already support multi-core horizontal scaling. With an extension to distributed horizontal scaling, they will be the ideal platform for diverse services, including those poorly served by Hadoop...

Akka for distributed computation.



Friday, March 9, 12

Akka is the right platform for distributed services. It exposes clean, low-level primitives for robust, distributed services (e.g., Actors), upon which we can build flexible big data systems that can handle soft real time and batch processing efficiently and scalably.
(No, this isn't Akka Mountain in Sweden. So sue me... ;)



Functional Programming

for Java Developers

O'REILLY®

Dean Wampler

FP for Java Devs
1/2 off today!!

30

Friday, March 9, 12

That's it. Today only (3/9), you can get my ebook 1/2 off!