

Motorola Solutions SABA
November 13, 2012
dean.wampler@thinkbiganalytics.com

MapReduce and its Discontents



Wednesday, November 14, 12

Is MapReduce the end of the “Big Data” story? Does it meet all our needs?

Big Data



It's a buzz word, but generally associated with the problem of data sets too big to manage with traditional SQL databases. A parallel development has been the NoSQL movement that is good at handling semistructured data, scaling, etc.

Big Data

Data so big that traditional solutions are too slow, too small, or too expensive to use.



It's a buzz word, but generally associated with the problem of data sets too big to manage with traditional SQL databases. A parallel development has been the NoSQL movement that is good at handling semistructured data, scaling, etc.



3

Wednesday, November 14, 12

Three trends influence my thinking...



3 Trends

3

Wednesday, November 14, 12

Three trends influence my thinking...

Data Size ↑

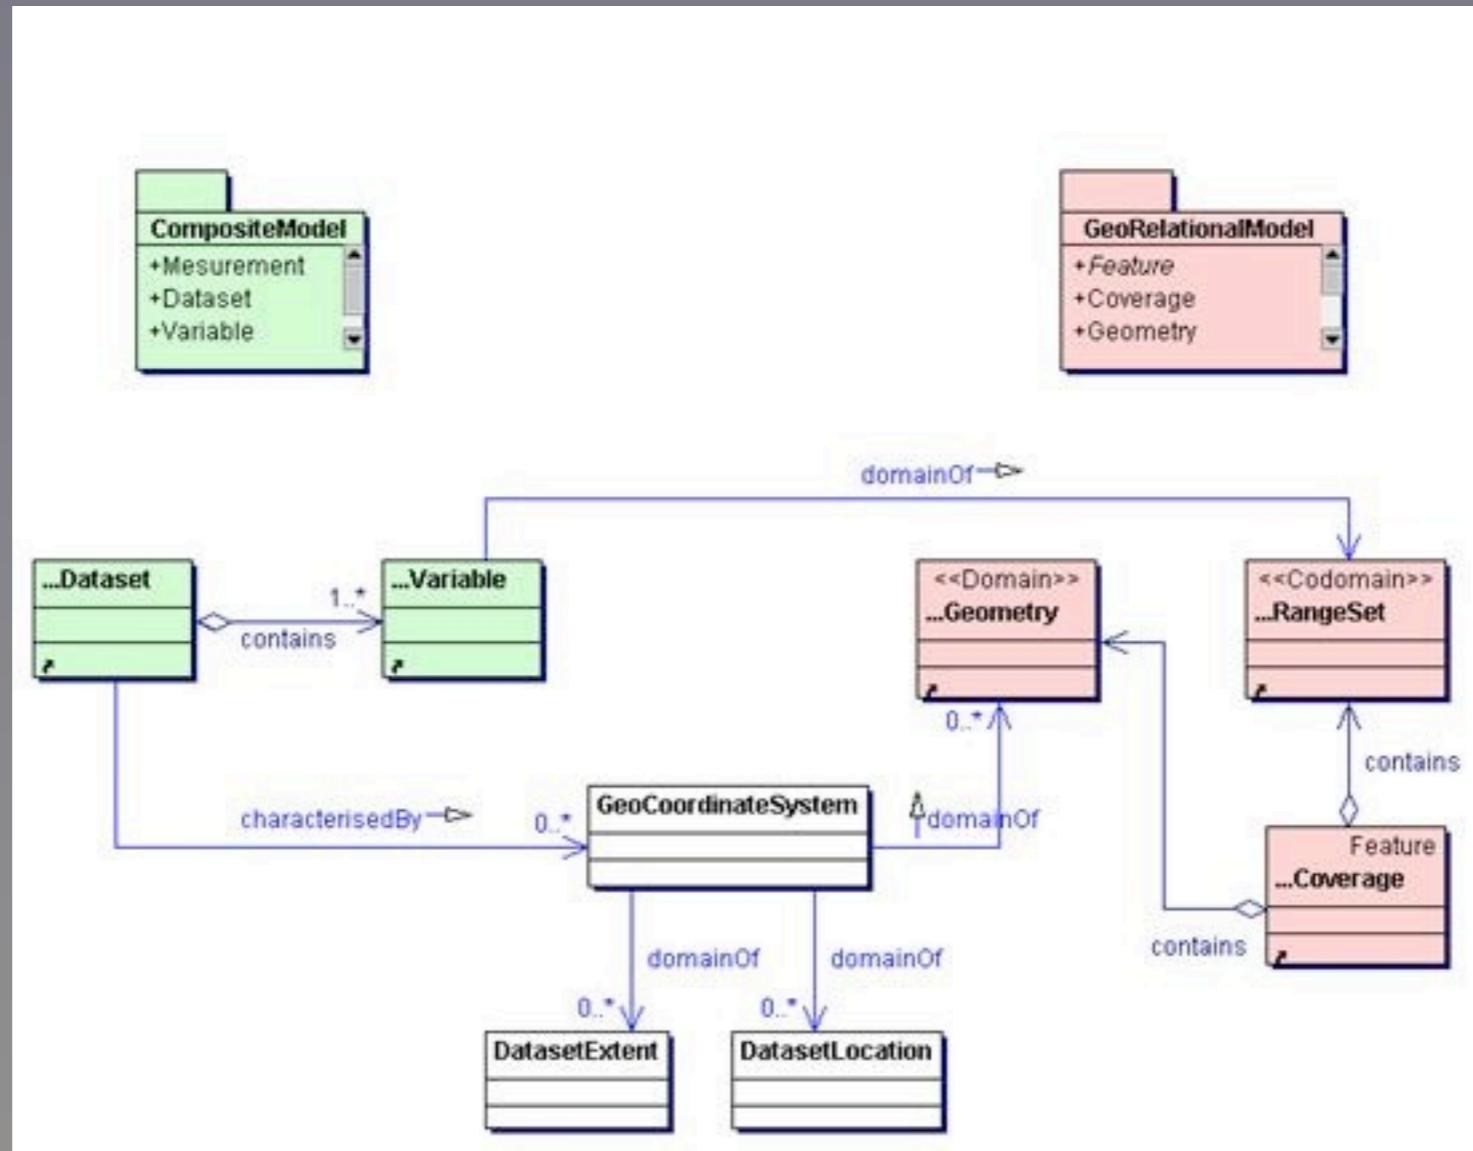


4

Wednesday, November 14, 12

Data volumes are obviously growing... rapidly.
Facebook now has over 600PB (Petabytes) of data!

Formal Schemas



5

Wednesday, November 14, 12

There is less emphasis on “formal” schemas and domain models, i.e., both relational models of data and OO models, because data schemas and sources change rapidly. So, using relatively-agnostic software (e.g., collections of things where the software is agnostic about the contents) tends to be faster to develop and run. Put another way, we find it more useful to build somewhat agnostic applications and drive their behavior through data.

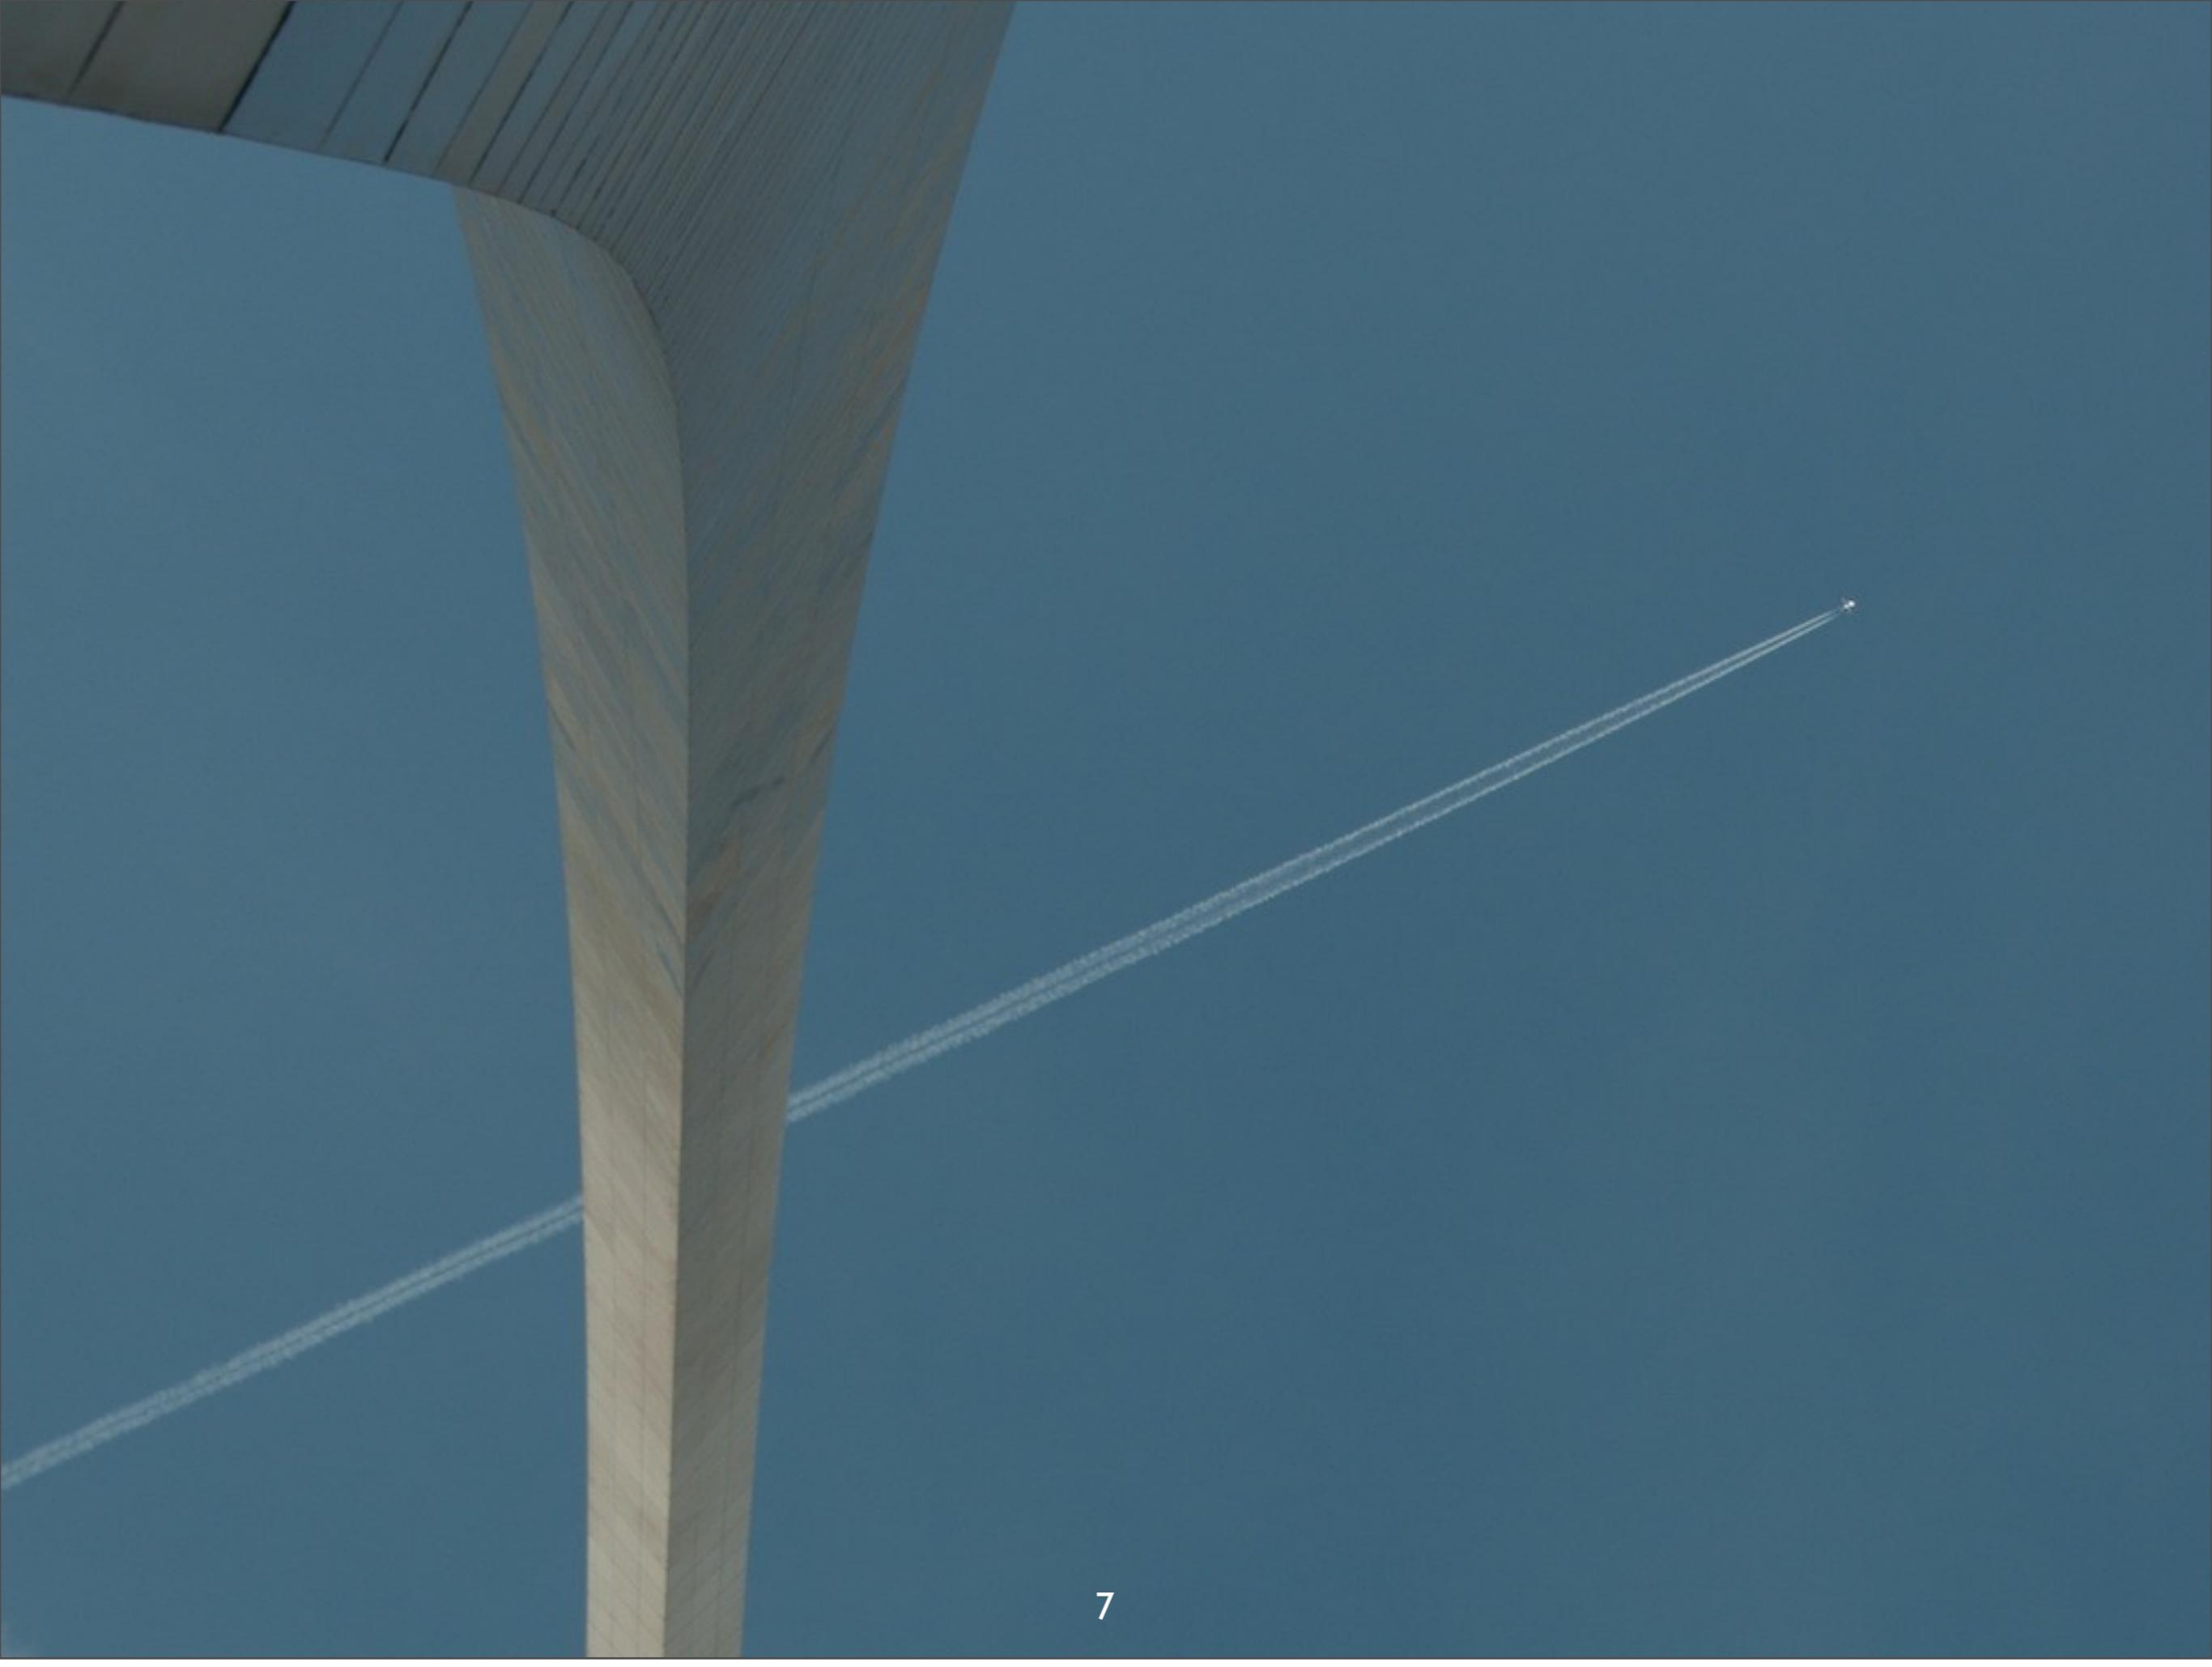
Data-Driven Programs ↑



6

Wednesday, November 14, 12

This is the 2nd generation “Stanley”, the most successful self-driving car ever built (by a Google-Stanford) team. Machine learning is growing in importance. Here, generic algorithms and data structures are trained to represent the “world” using data, rather than encoding a model of the world in the software itself. It’s another example of generic algorithms that produce the desired behavior by being application agnostic and data driven, rather than hard-coding a model of the world. (In practice, however, a balance is struck between completely agnostic apps and some engineering towards the specific problem.)



7

Wednesday, November 14, 12

What should software architectures look like for these kinds of systems?

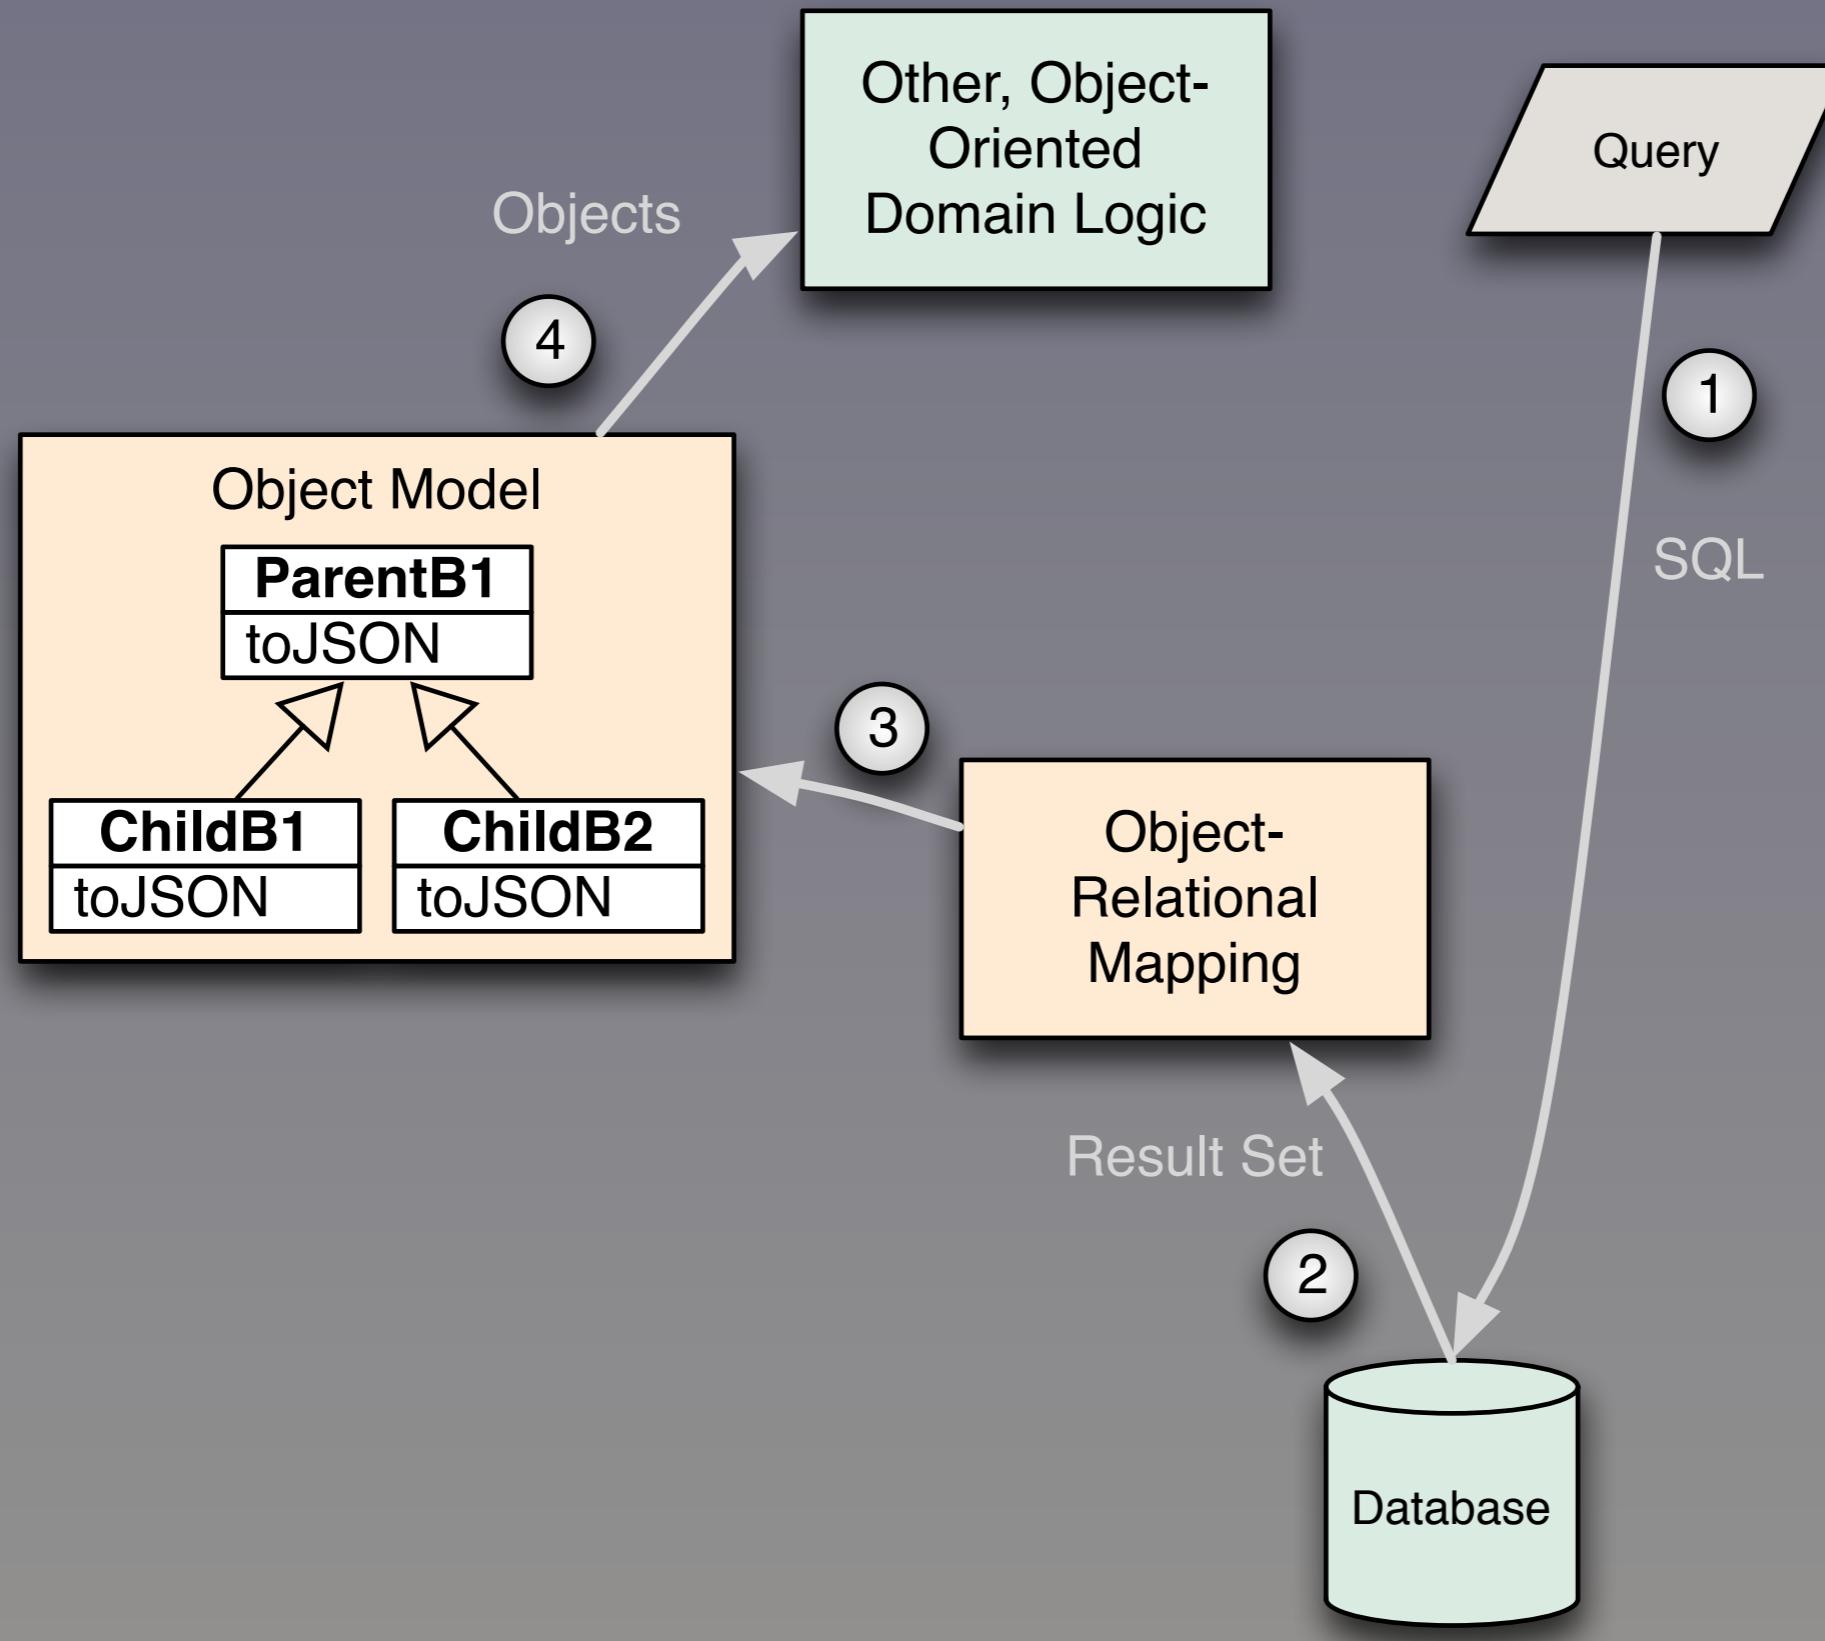


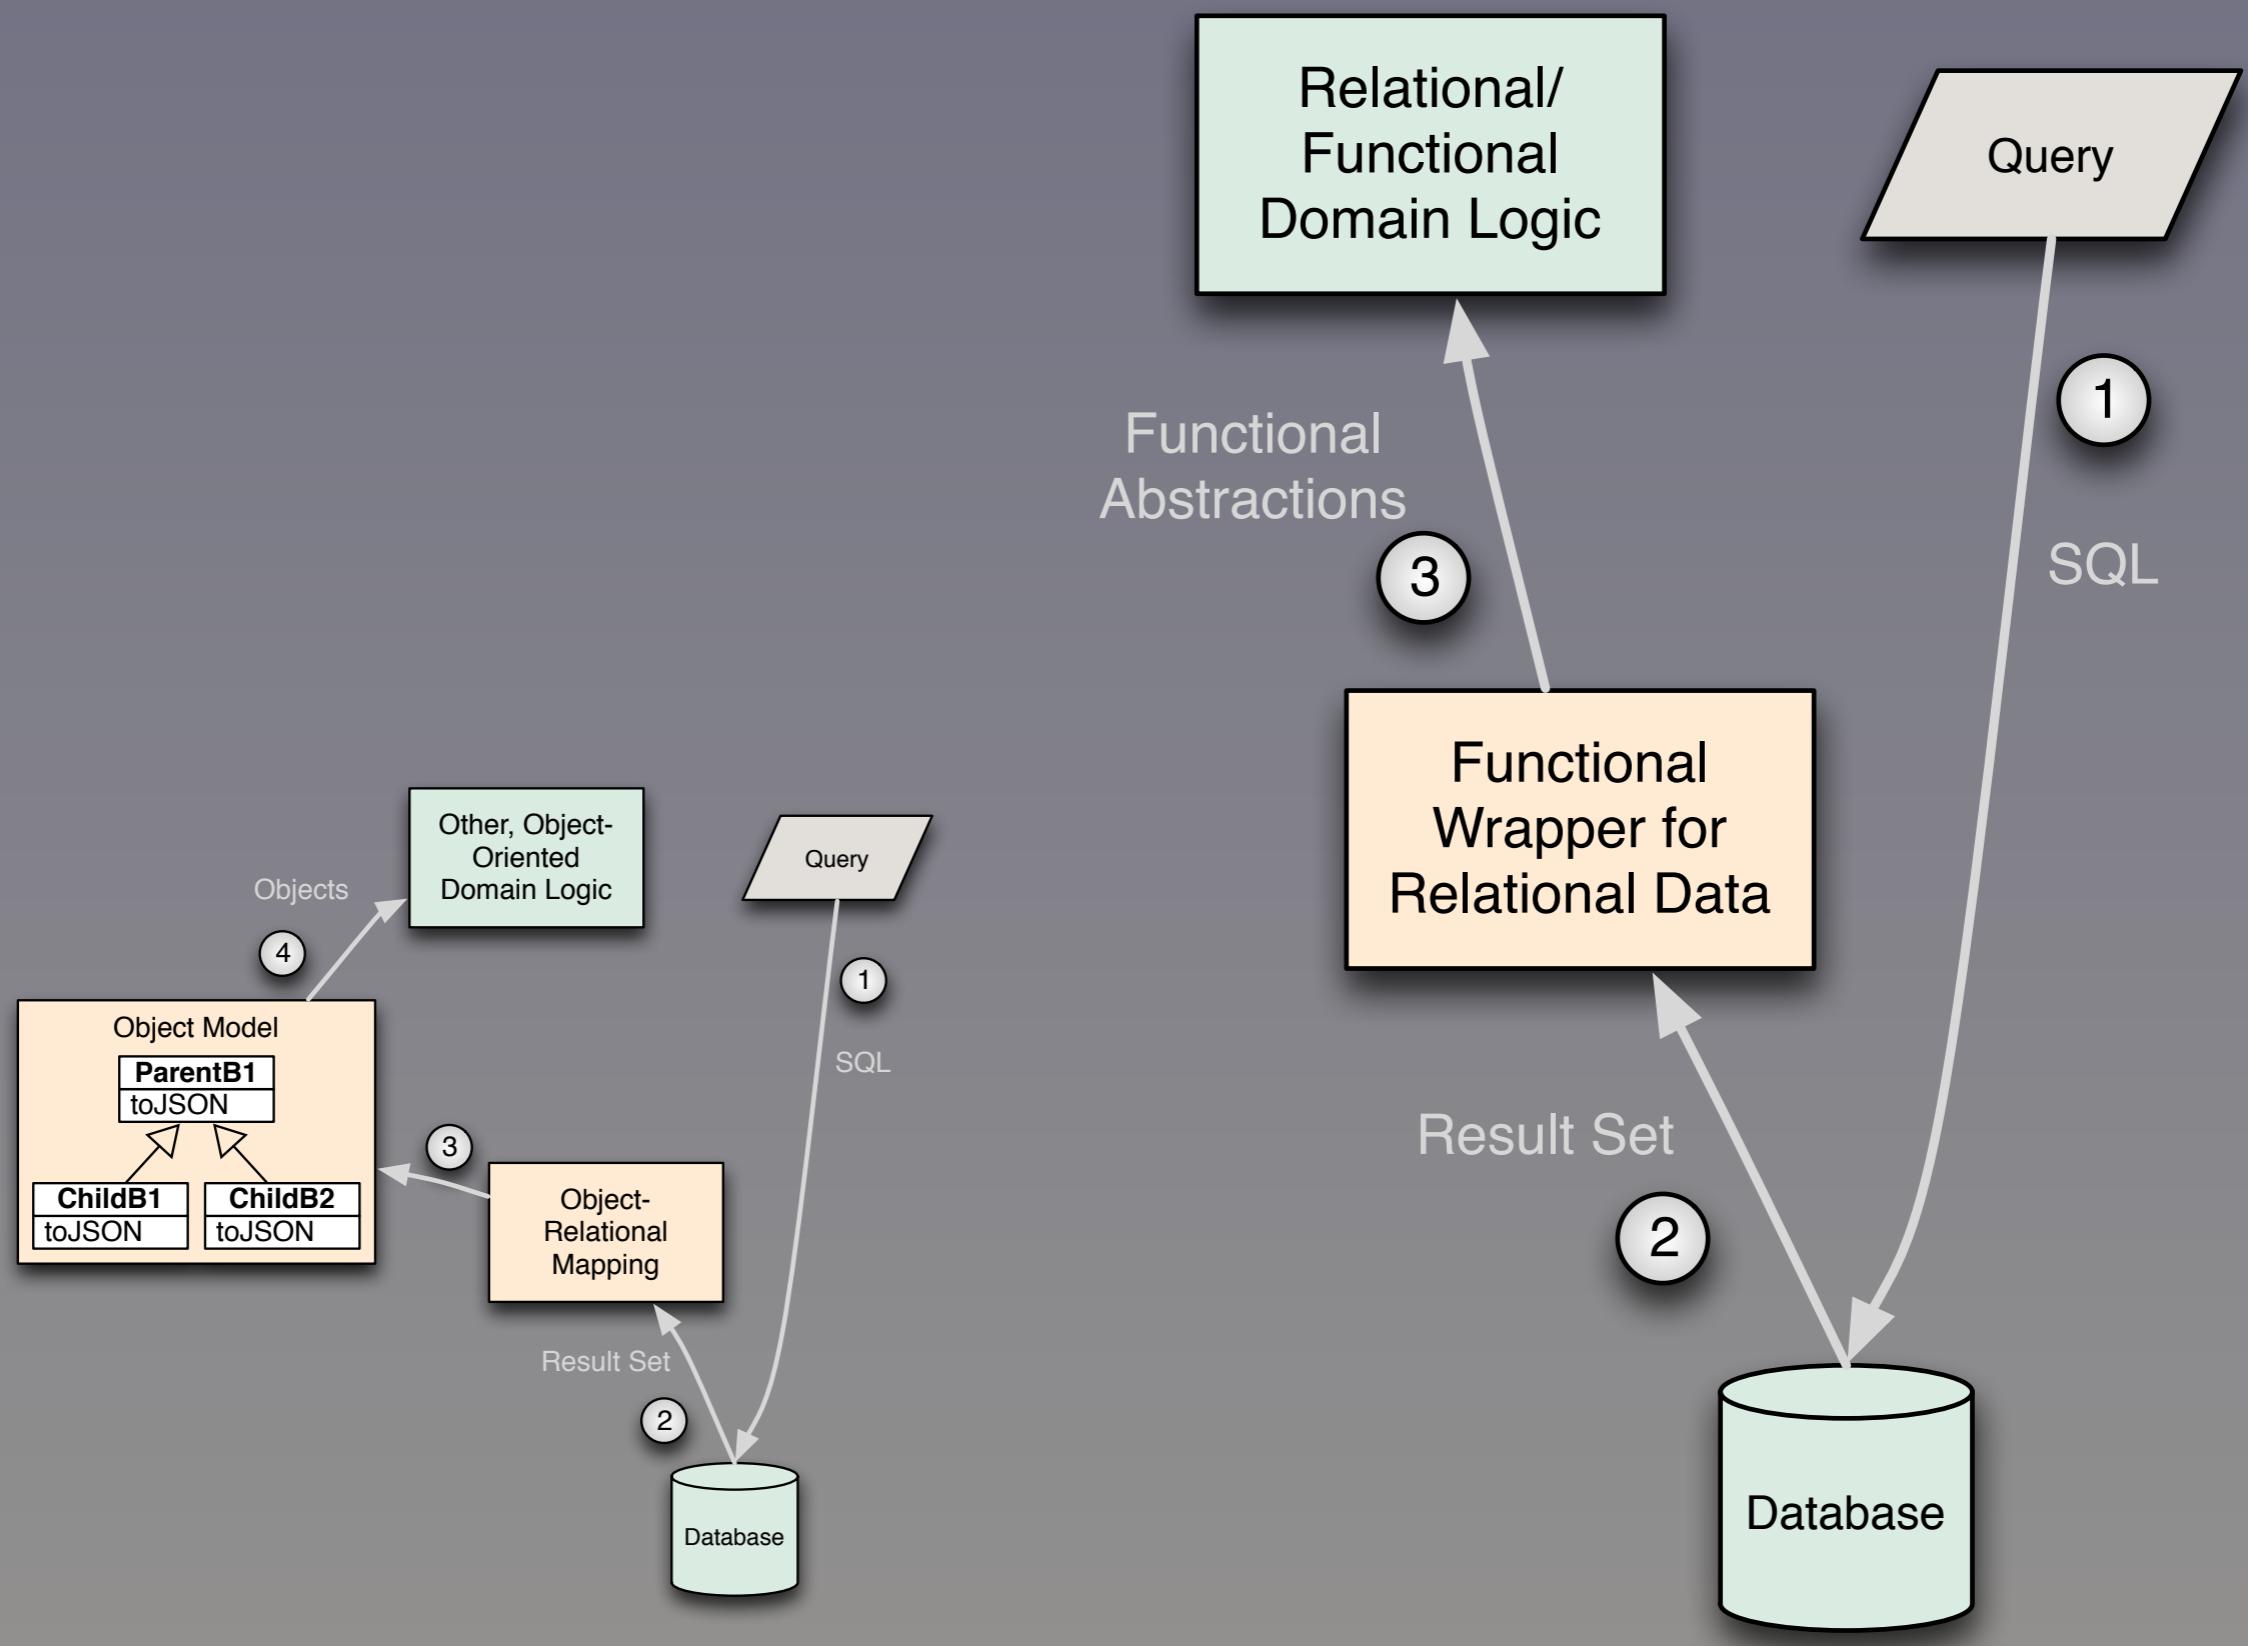
Big Data Architectures

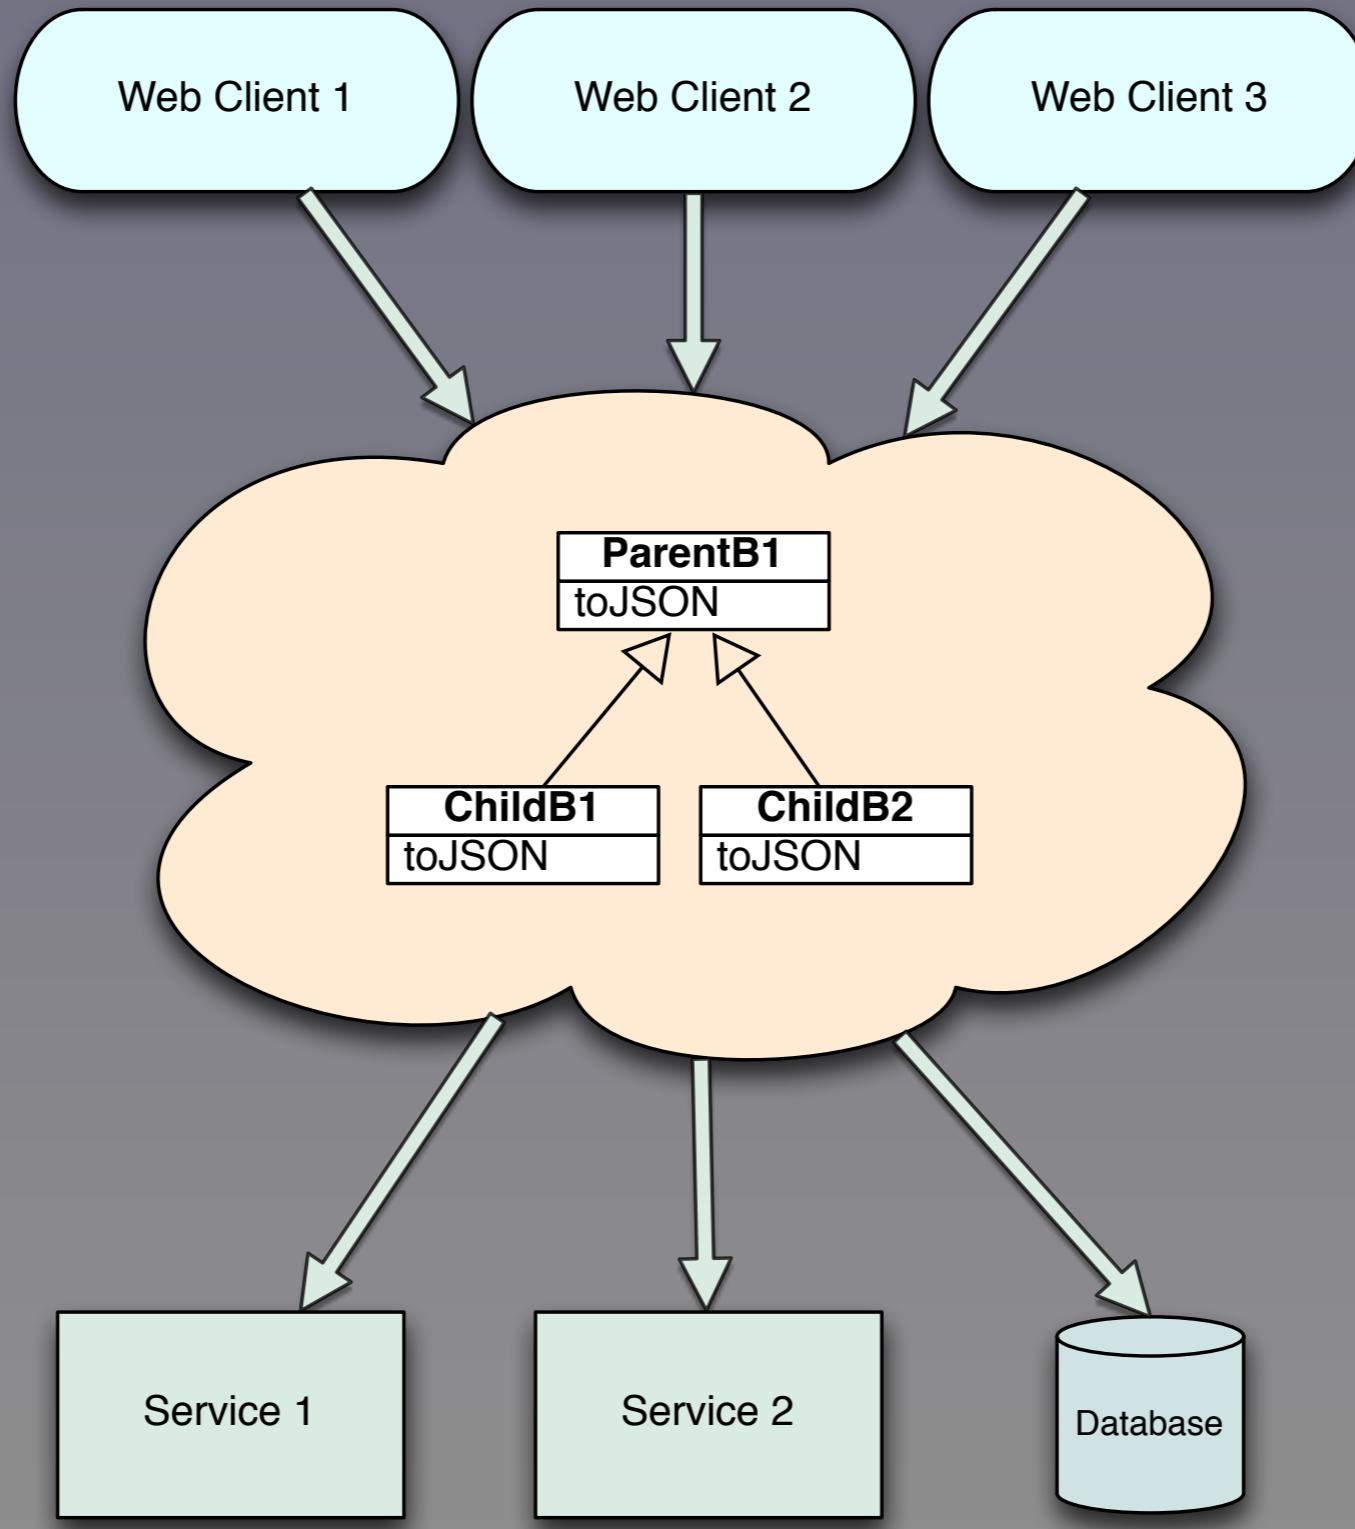
7

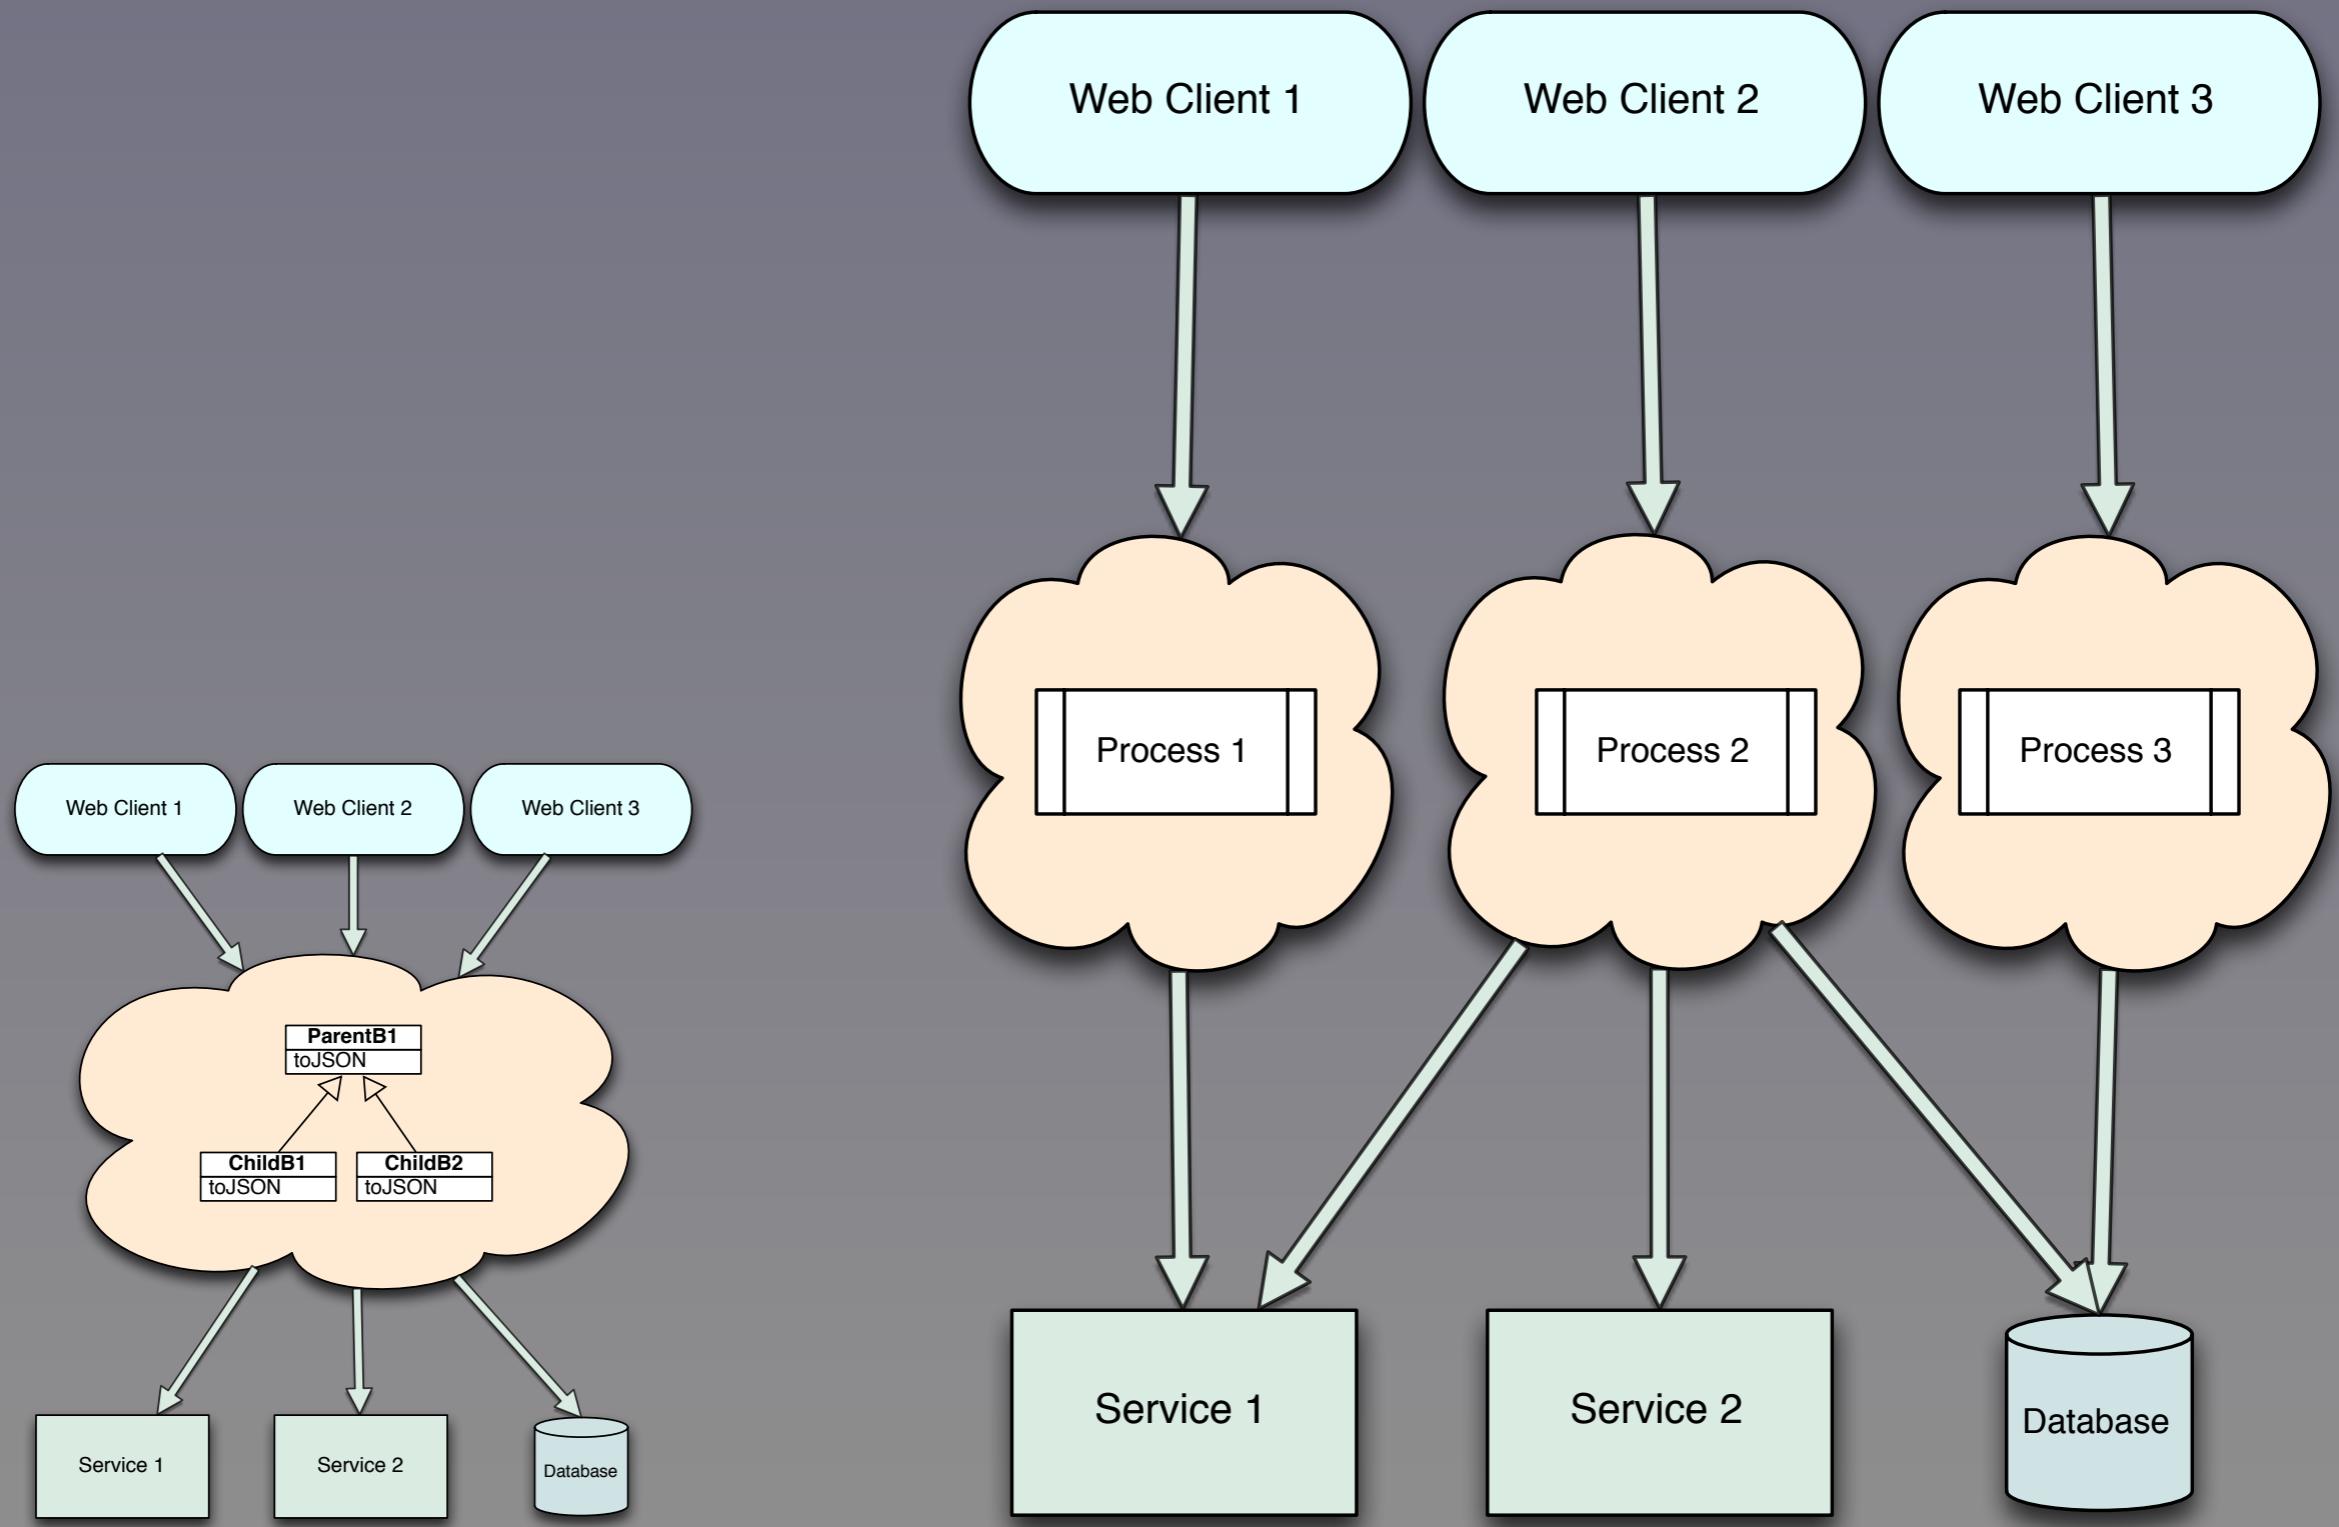
Wednesday, November 14, 12

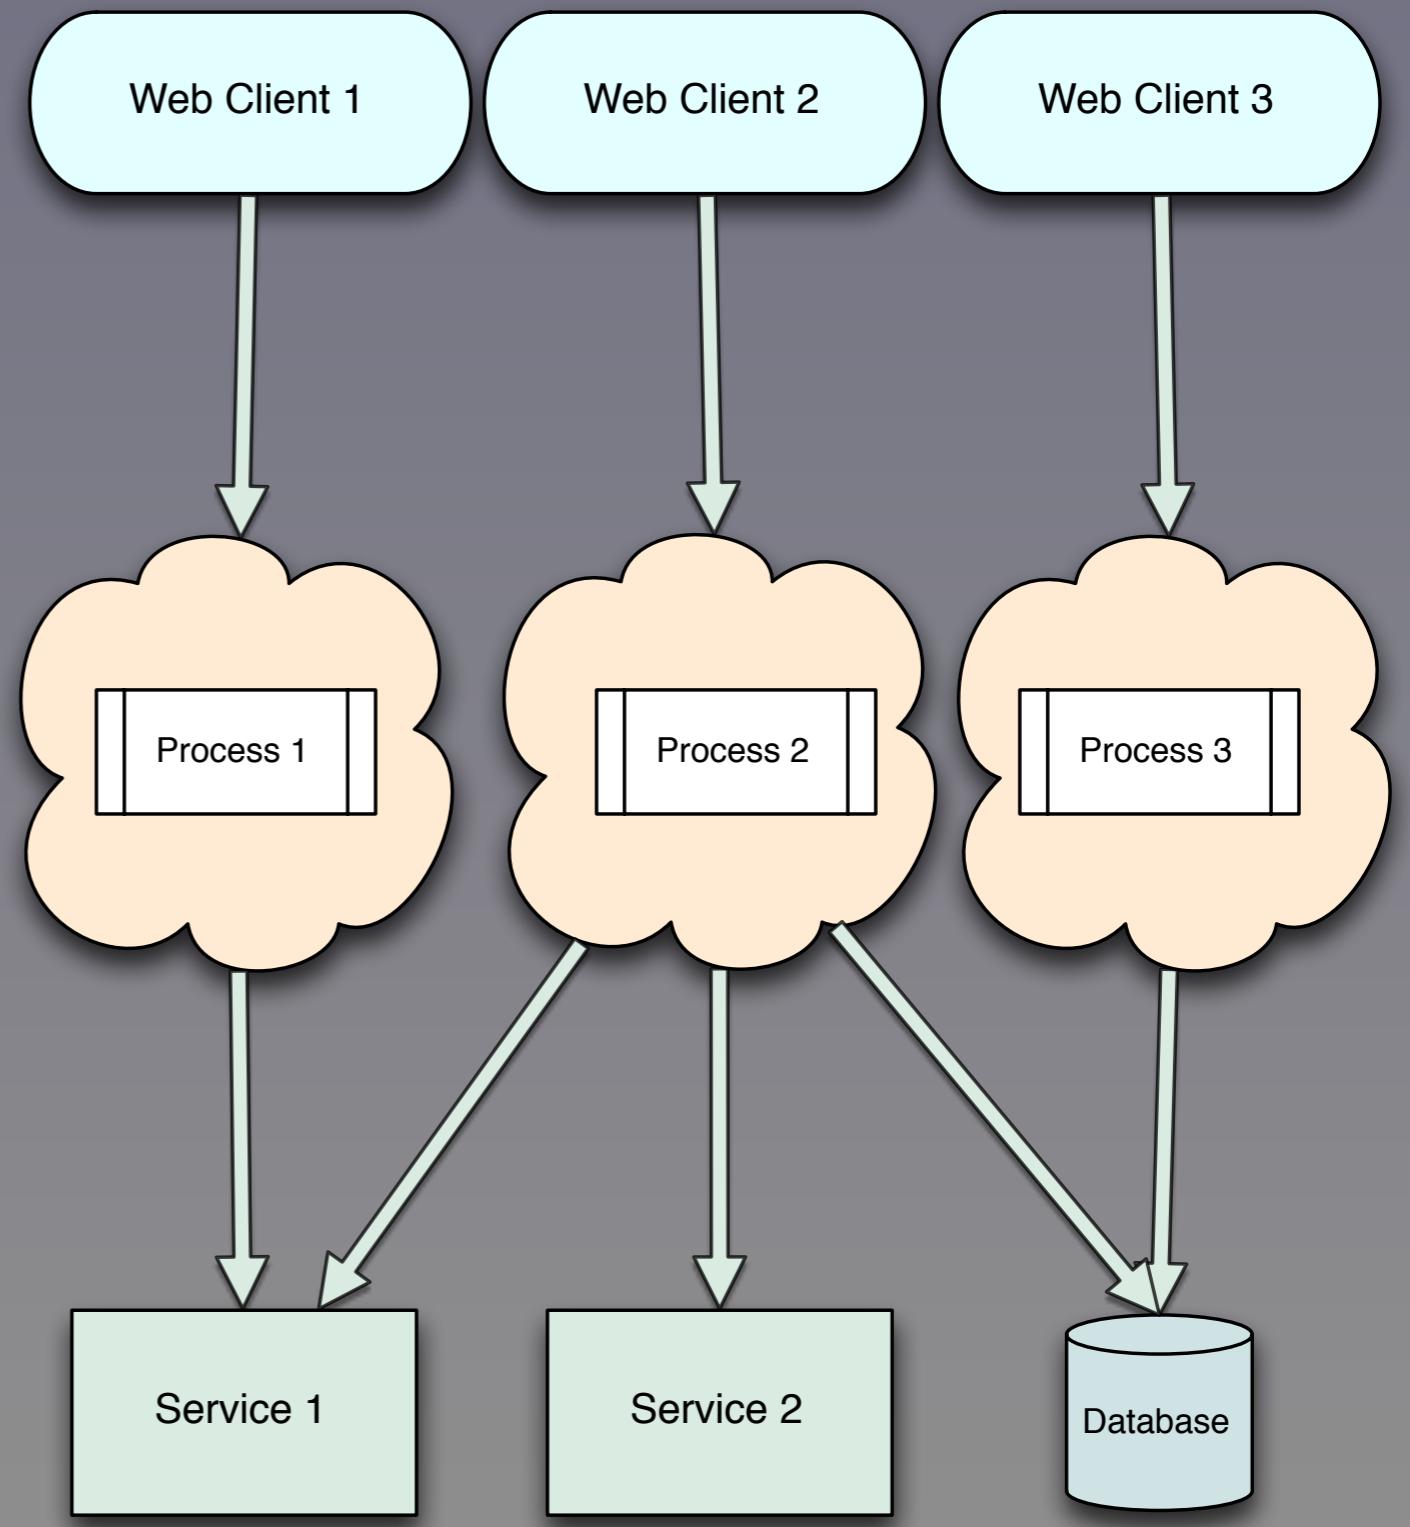
What should software architectures look like for these kinds of systems?



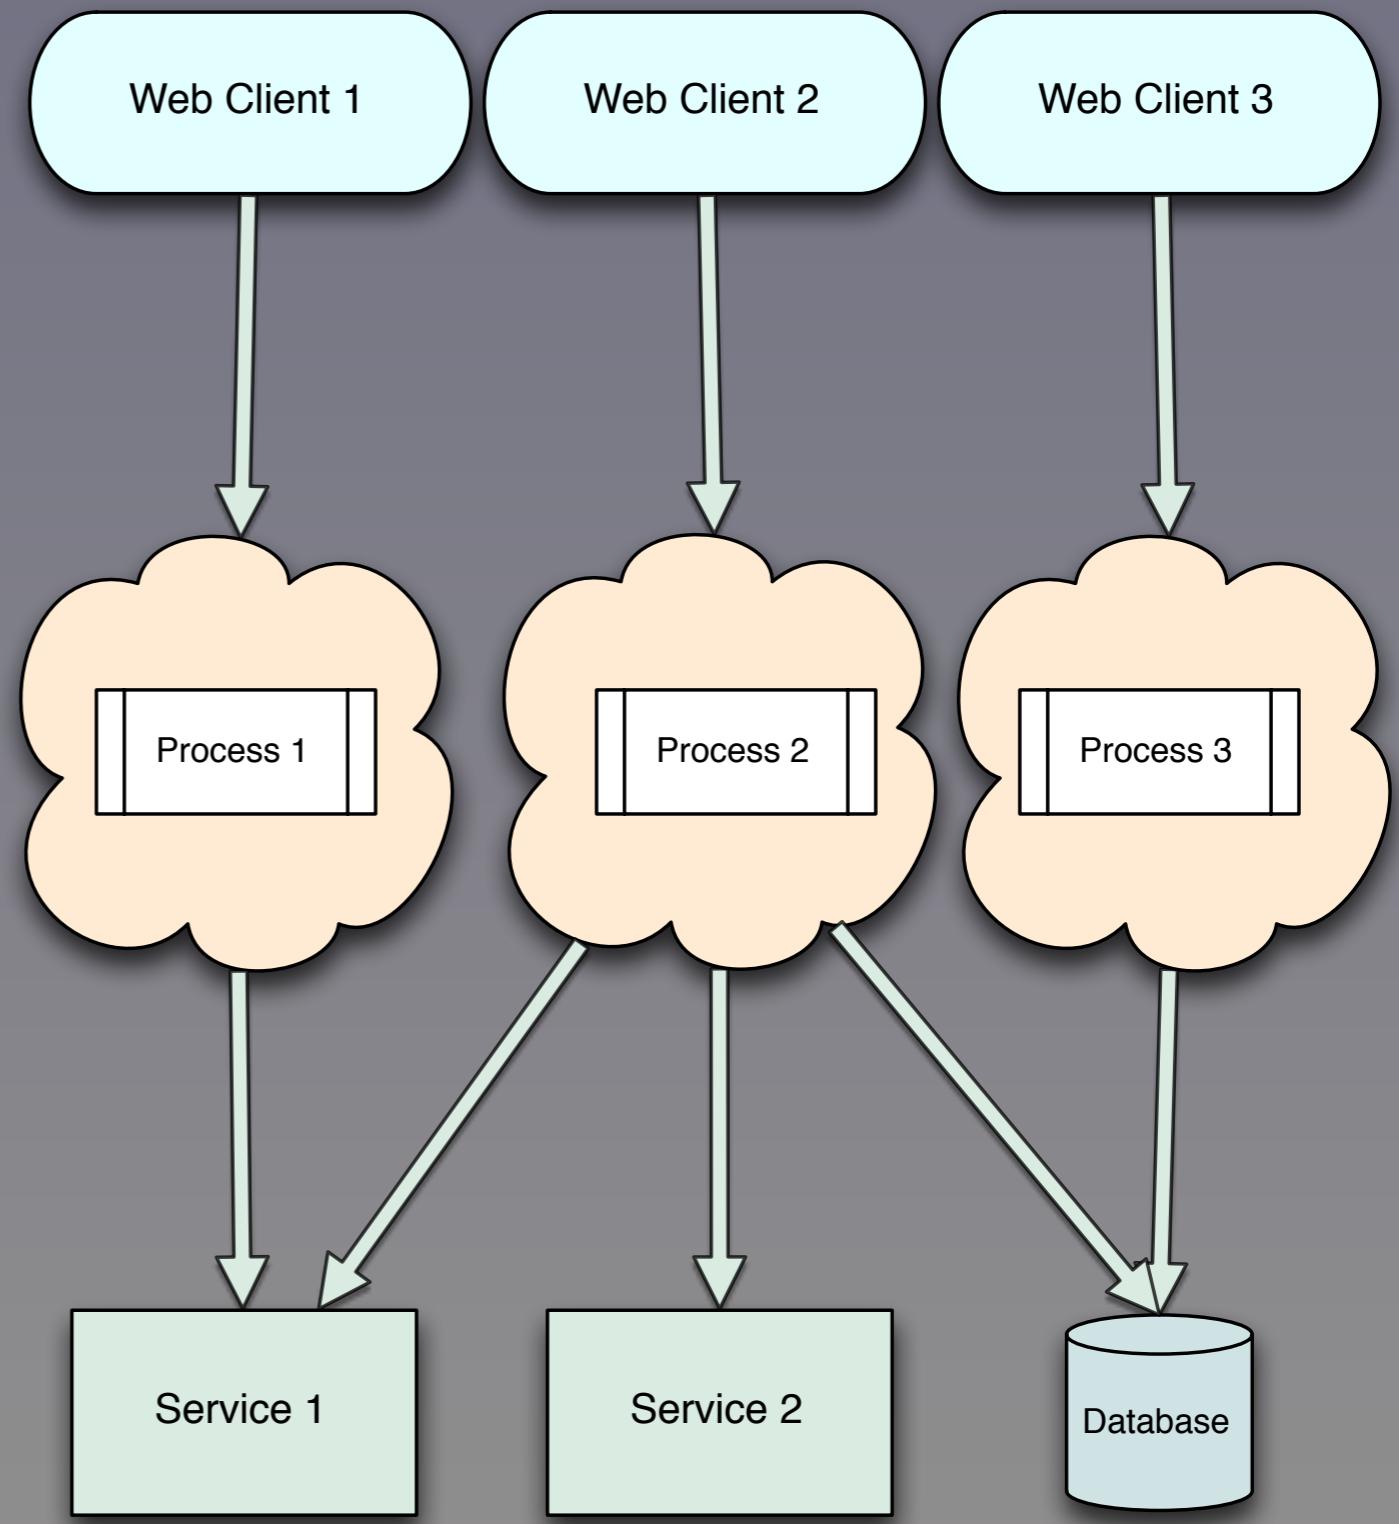


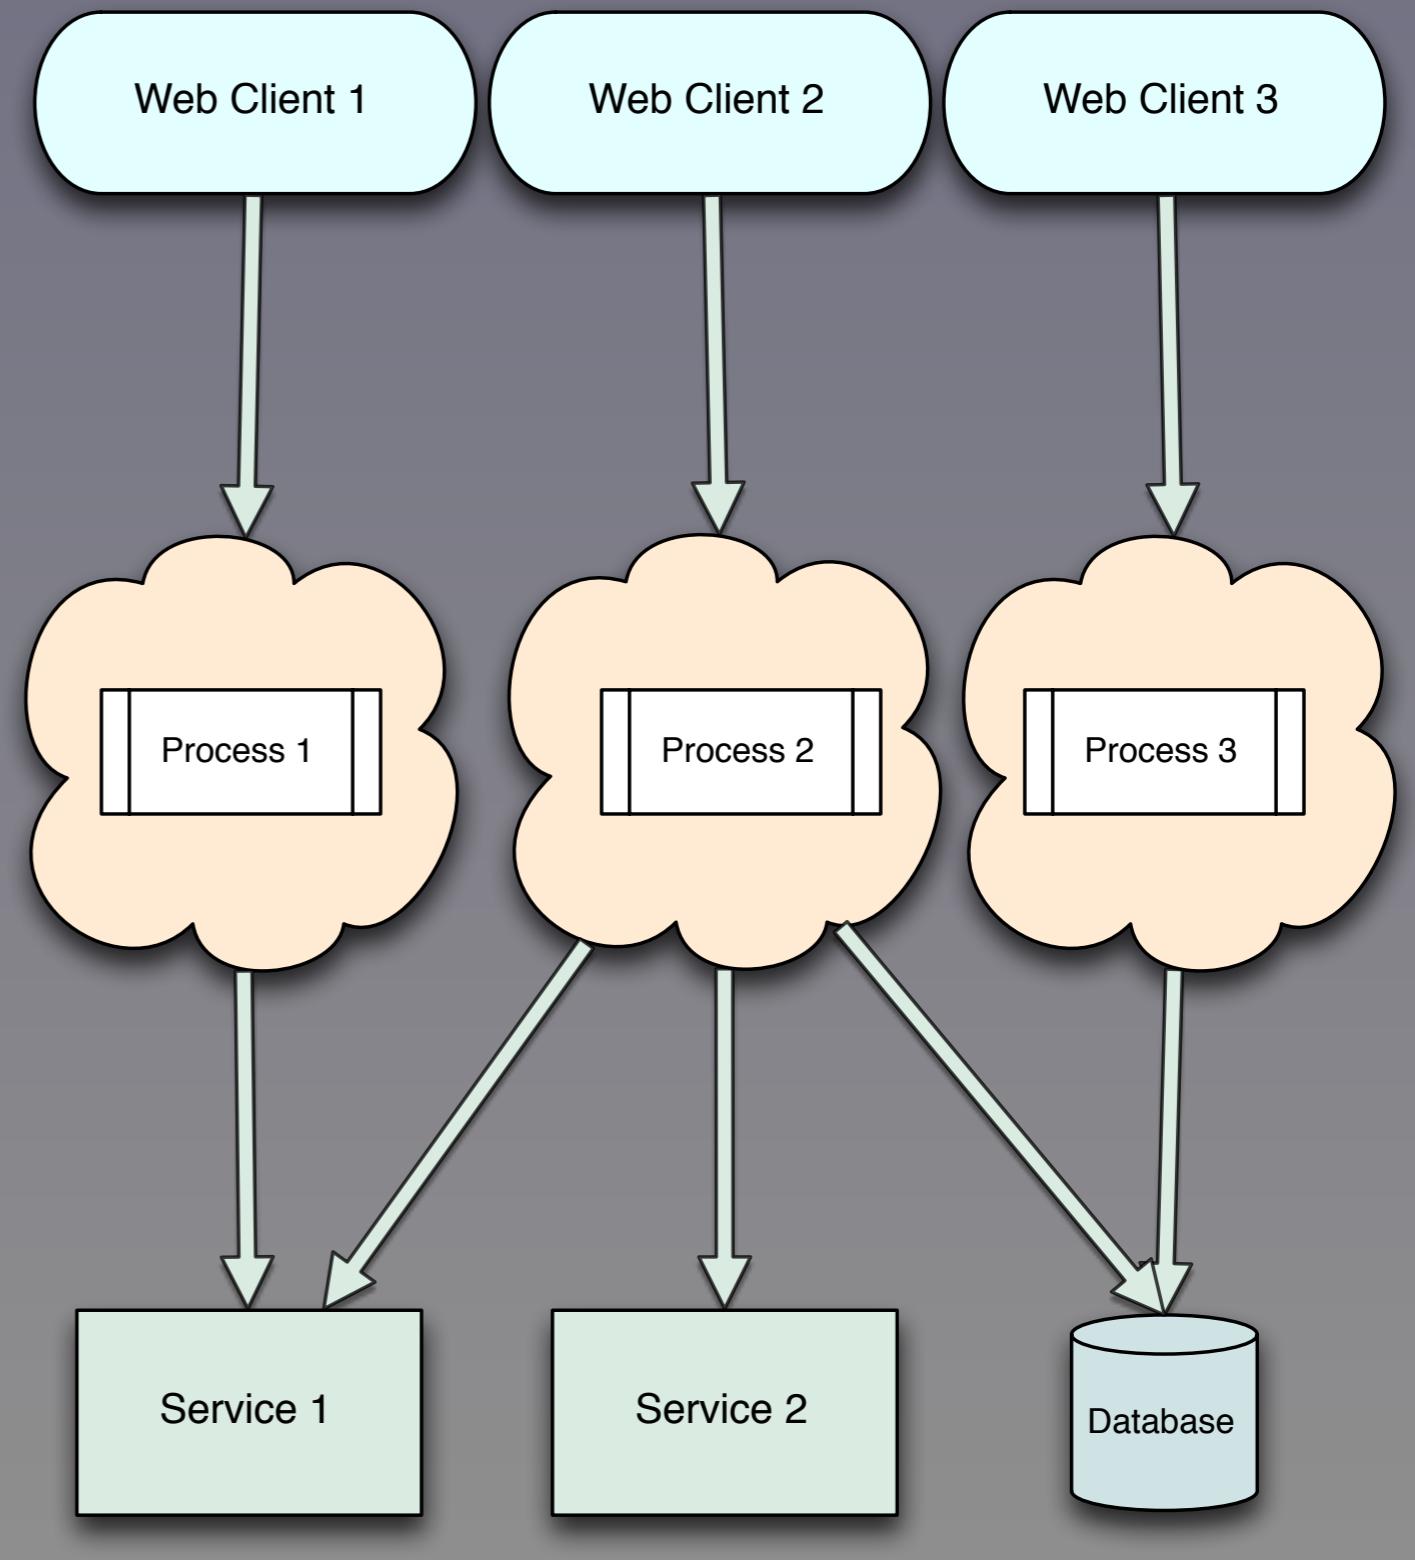






- Data Size ↑
- Formal Schema ↓
- Data-Driven Programs ↑



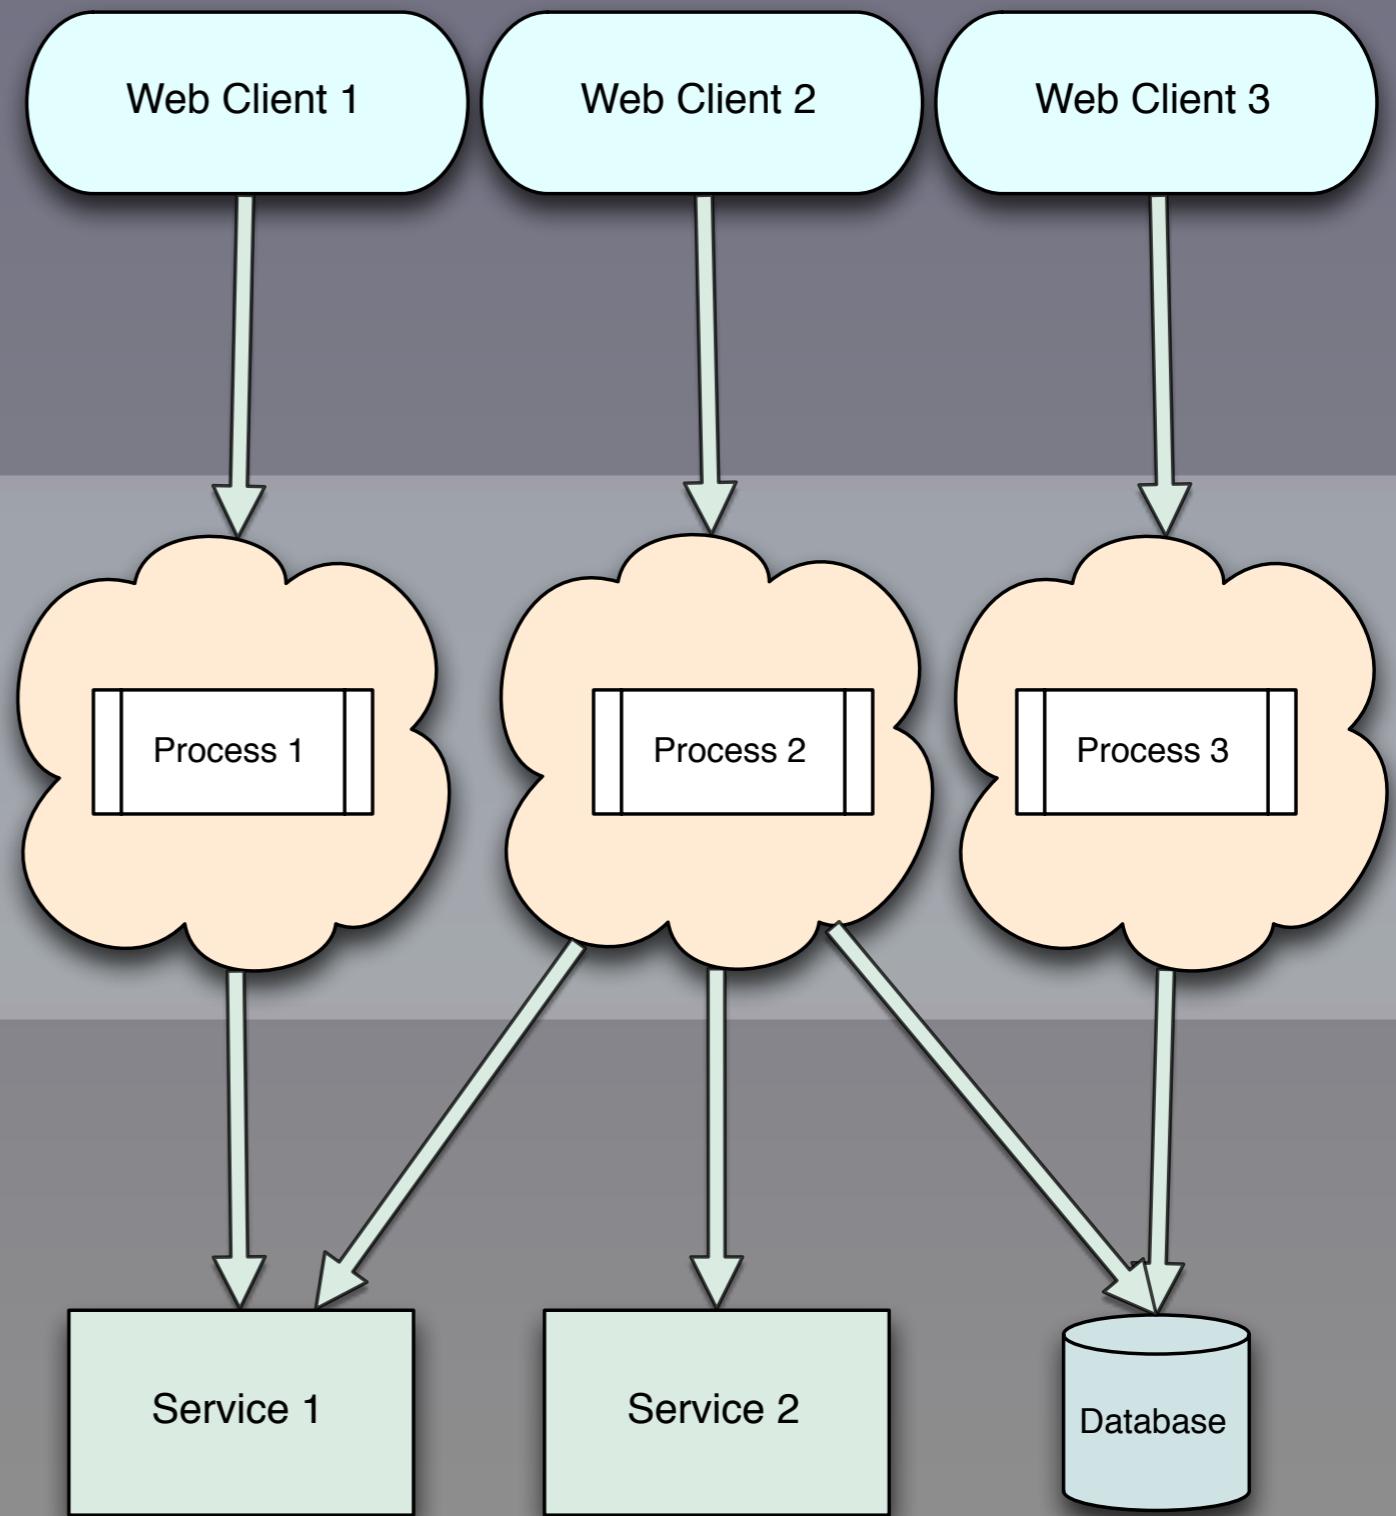


||

Wednesday, November 14, 12

And MapReduce + a distributed file system, like the Hadoop DFS, fit this model.

- MapReduce

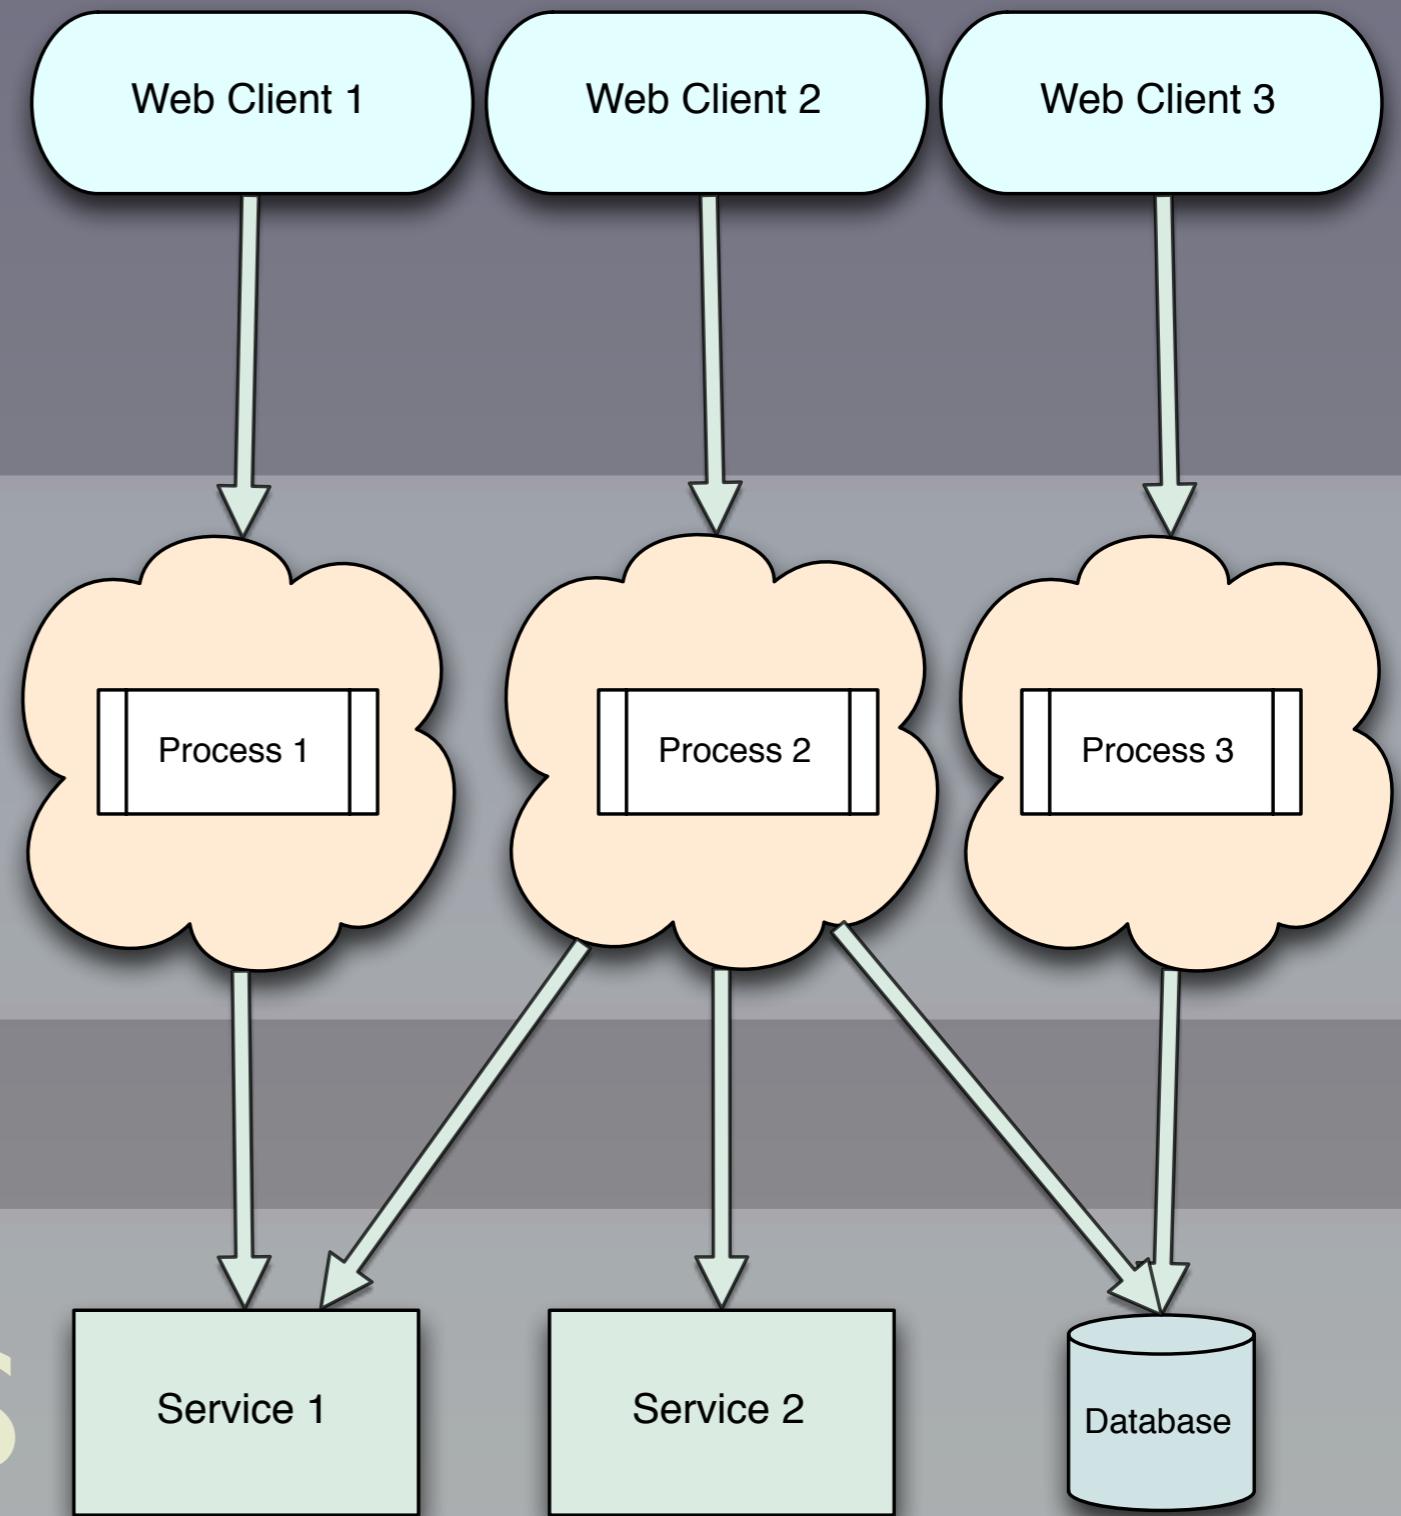


II

Wednesday, November 14, 12

And MapReduce + a distributed file system, like the Hadoop DFS, fit this model.

- MapReduce



- Distributed FS

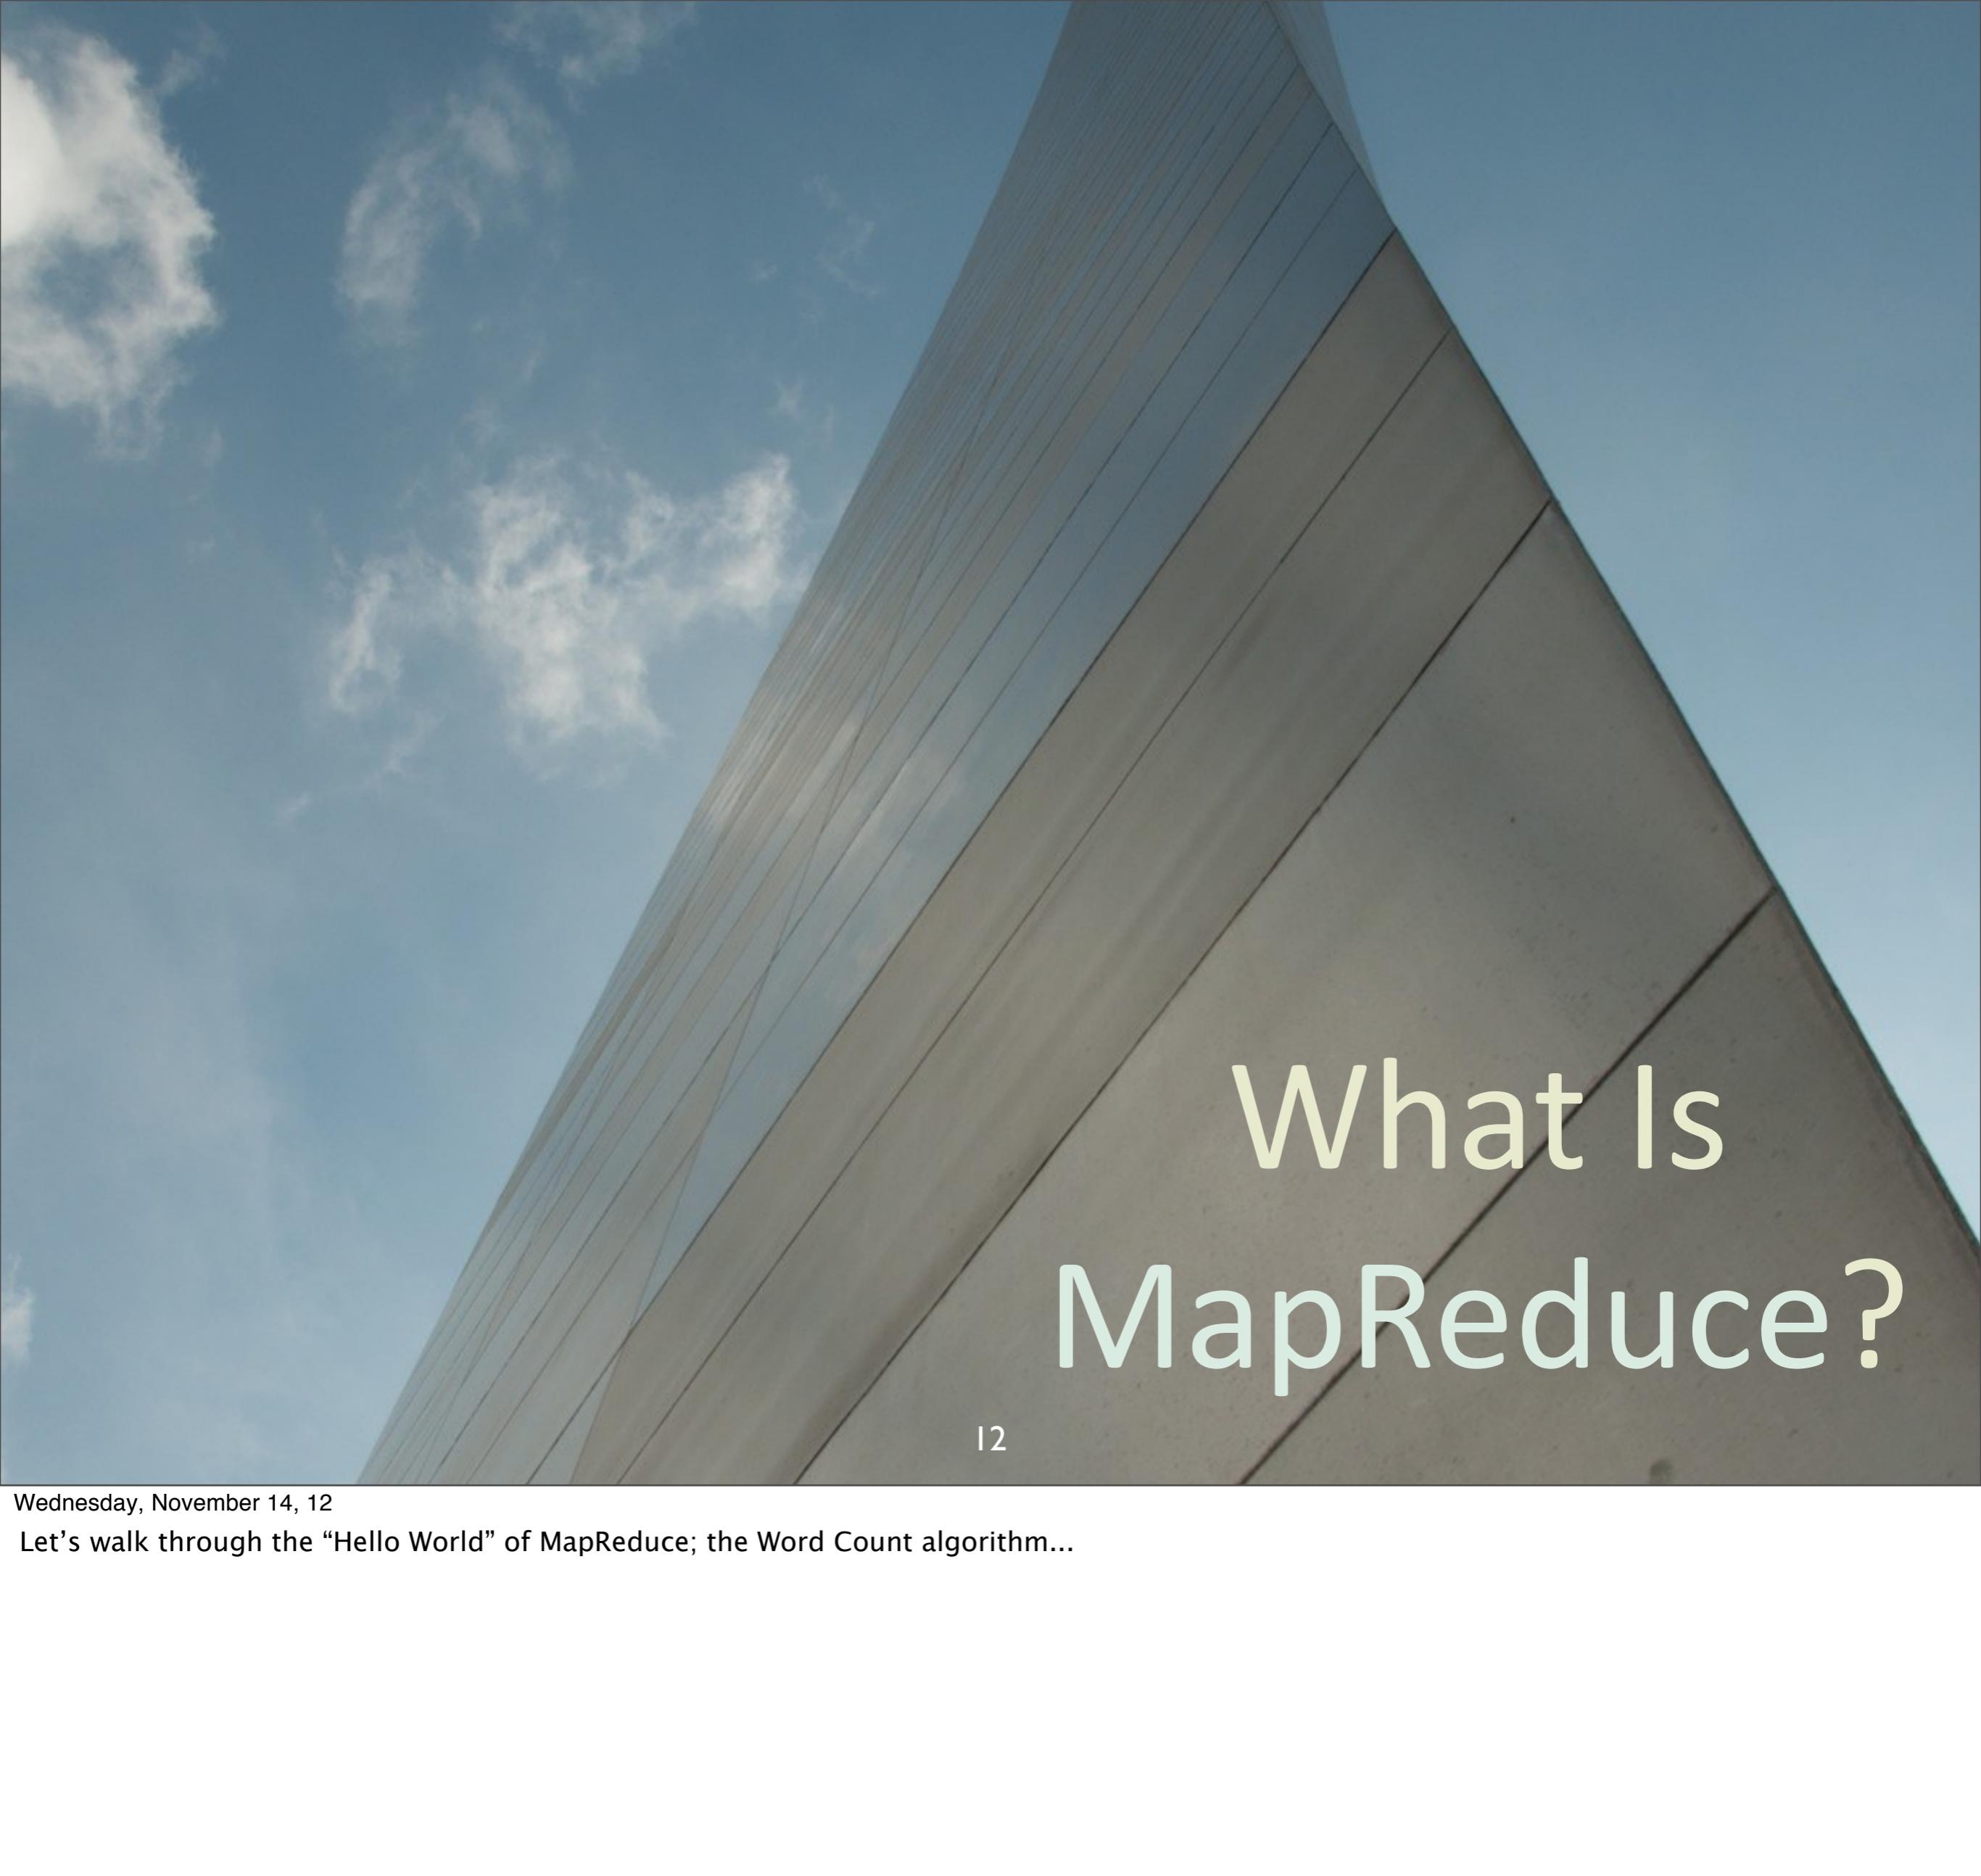
II

Wednesday, November 14, 12

And MapReduce + a distributed file system, like the Hadoop DFS, fit this model.

Wednesday, November 14, 12

Let's walk through the "Hello World" of MapReduce; the Word Count algorithm...

The background image shows a close-up view of a modern building's exterior. The building has a unique, curved design with a facade made of many thin, light-colored panels that create a textured, almost pleated appearance. The sky above is a clear, pale blue with some wispy white clouds.

What Is MapReduce?

12

Wednesday, November 14, 12

Let's walk through the "Hello World" of MapReduce; the Word Count algorithm...

MapReduce in Hadoop

Let's look at a
MapReduce algorithm:
WordCount.

Wednesday, November 14, 12

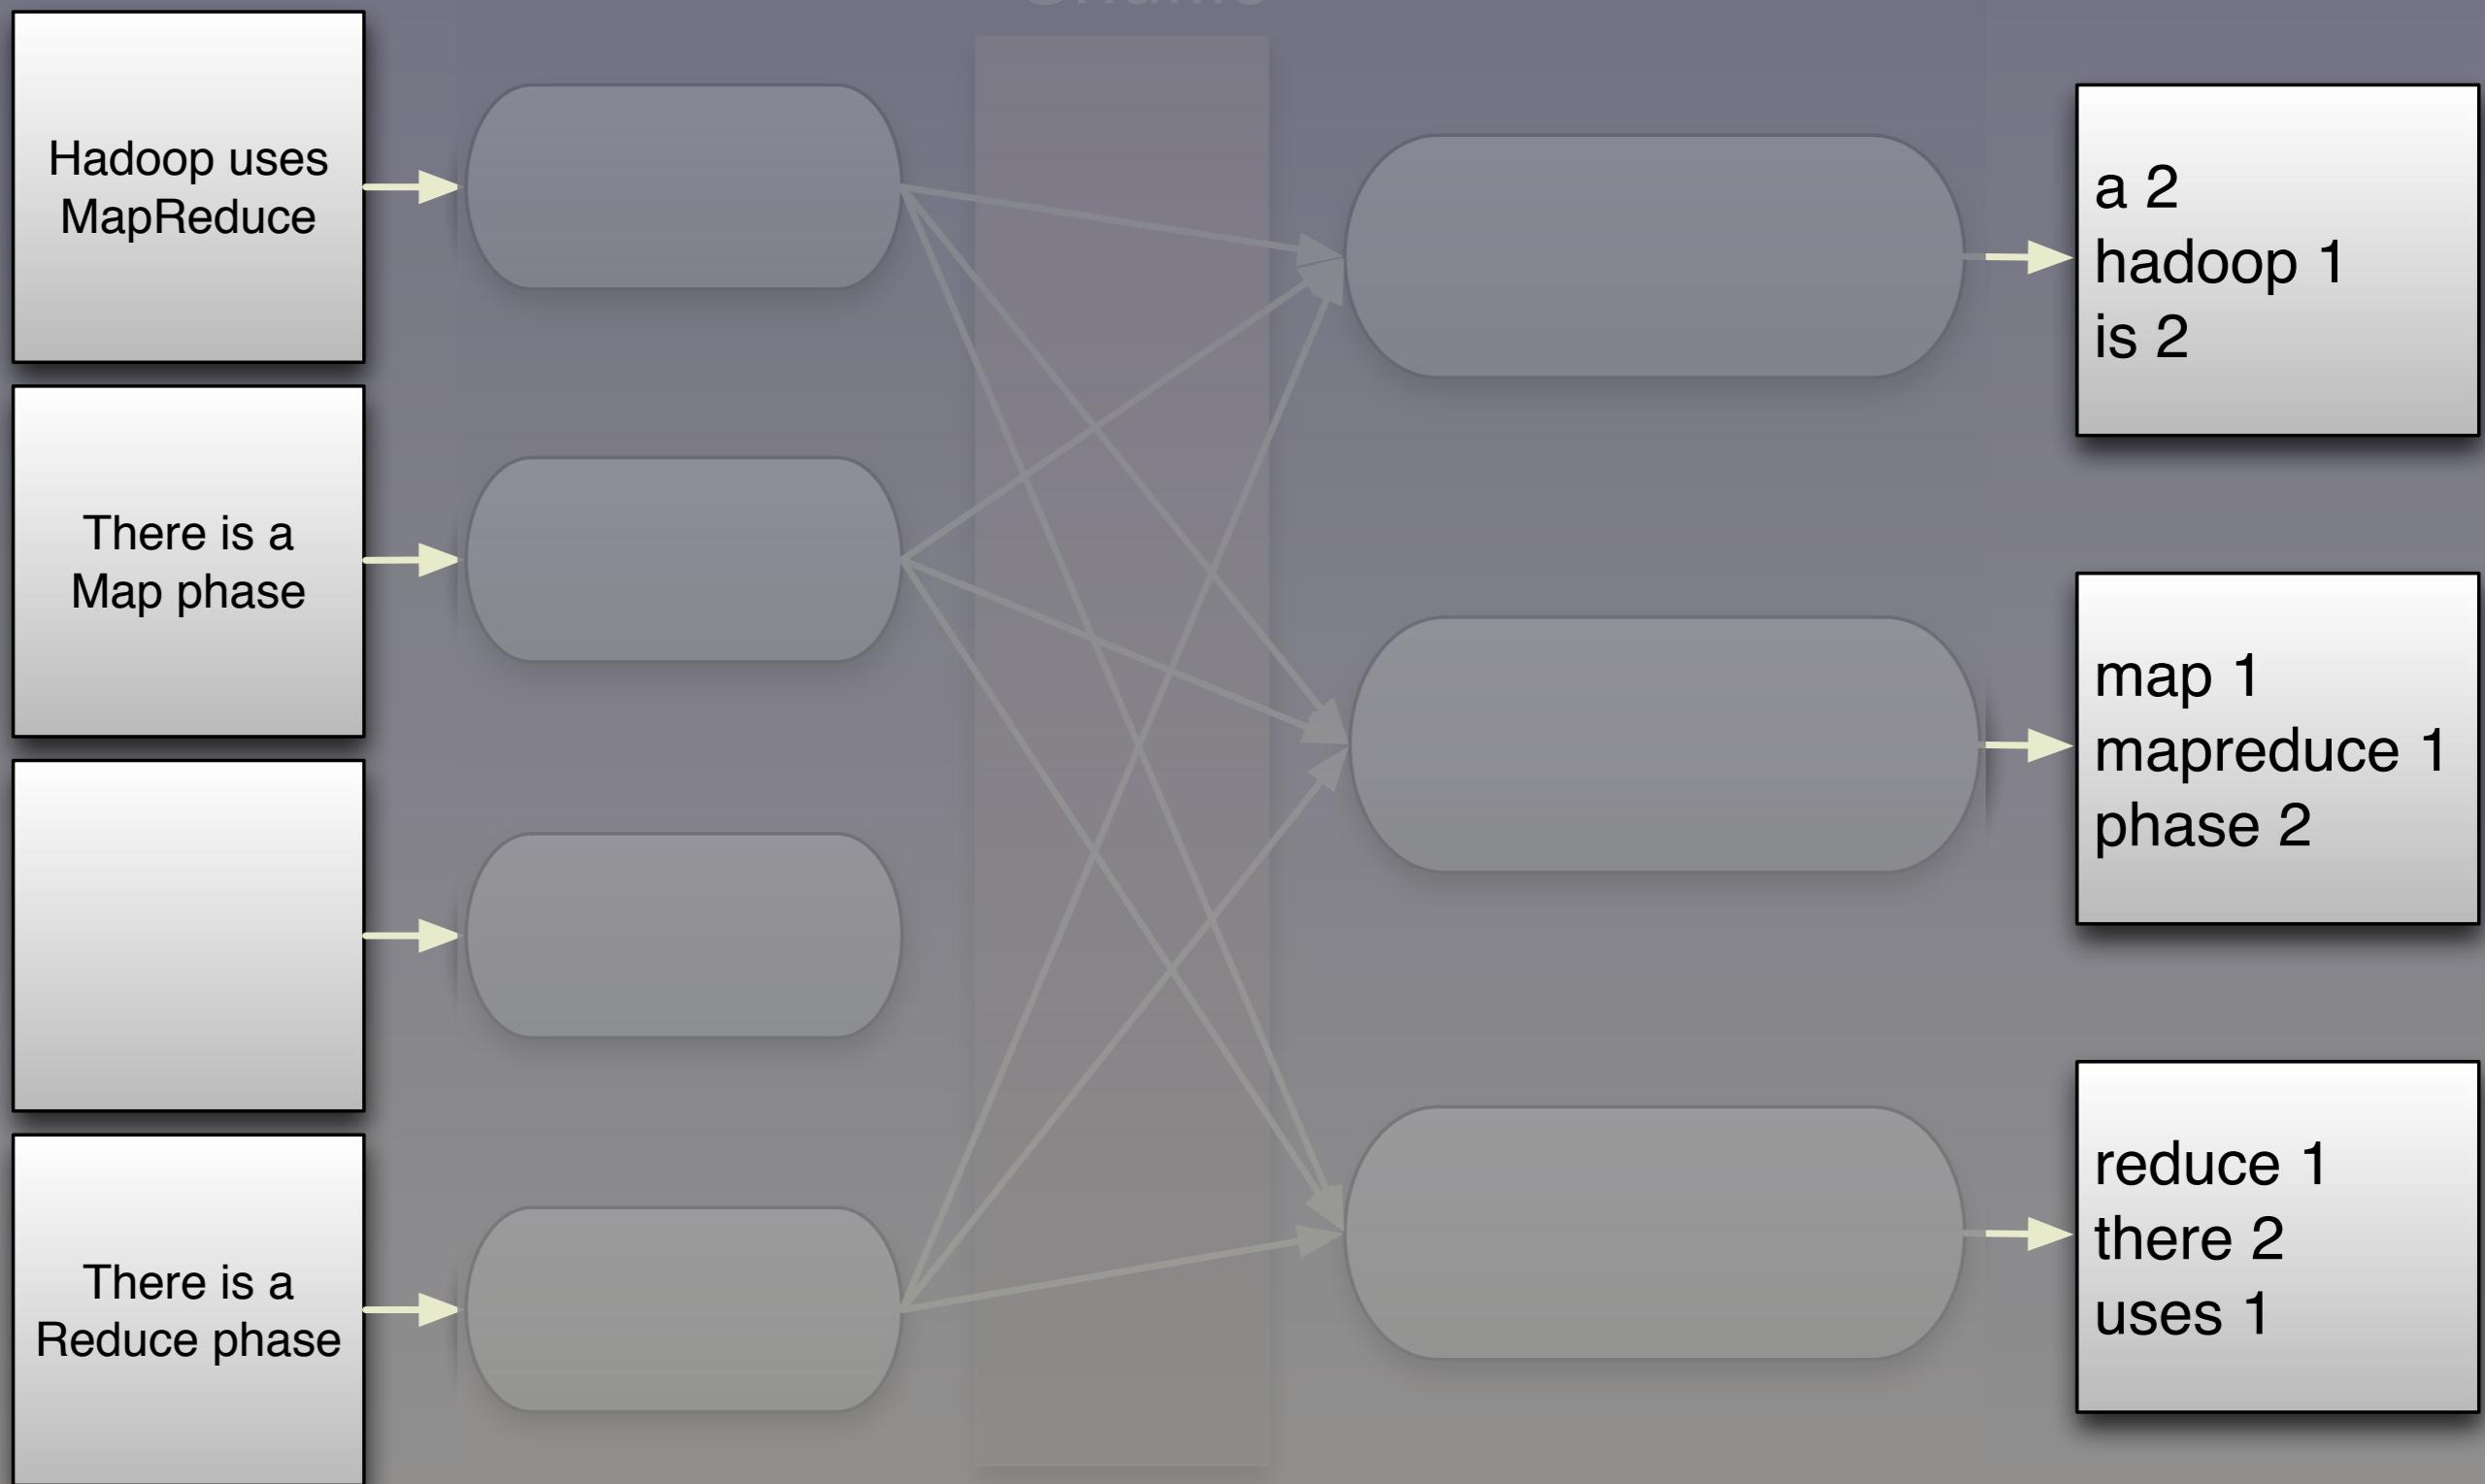
Let's walk through this algorithm at a conceptual level...

MapReduce in Hadoop

Let's look at a
MapReduce algorithm:
WordCount.

(The *Hello World* of big data...)

Input Mappers Sort, Shuffle Reducers Output



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Four input documents, one left empty, the others with small phrases (for simplicity...). The word count output is on the right (we'll see why there are three output "documents"). We need to get from the input on the left-hand side to the output on the right-hand side.

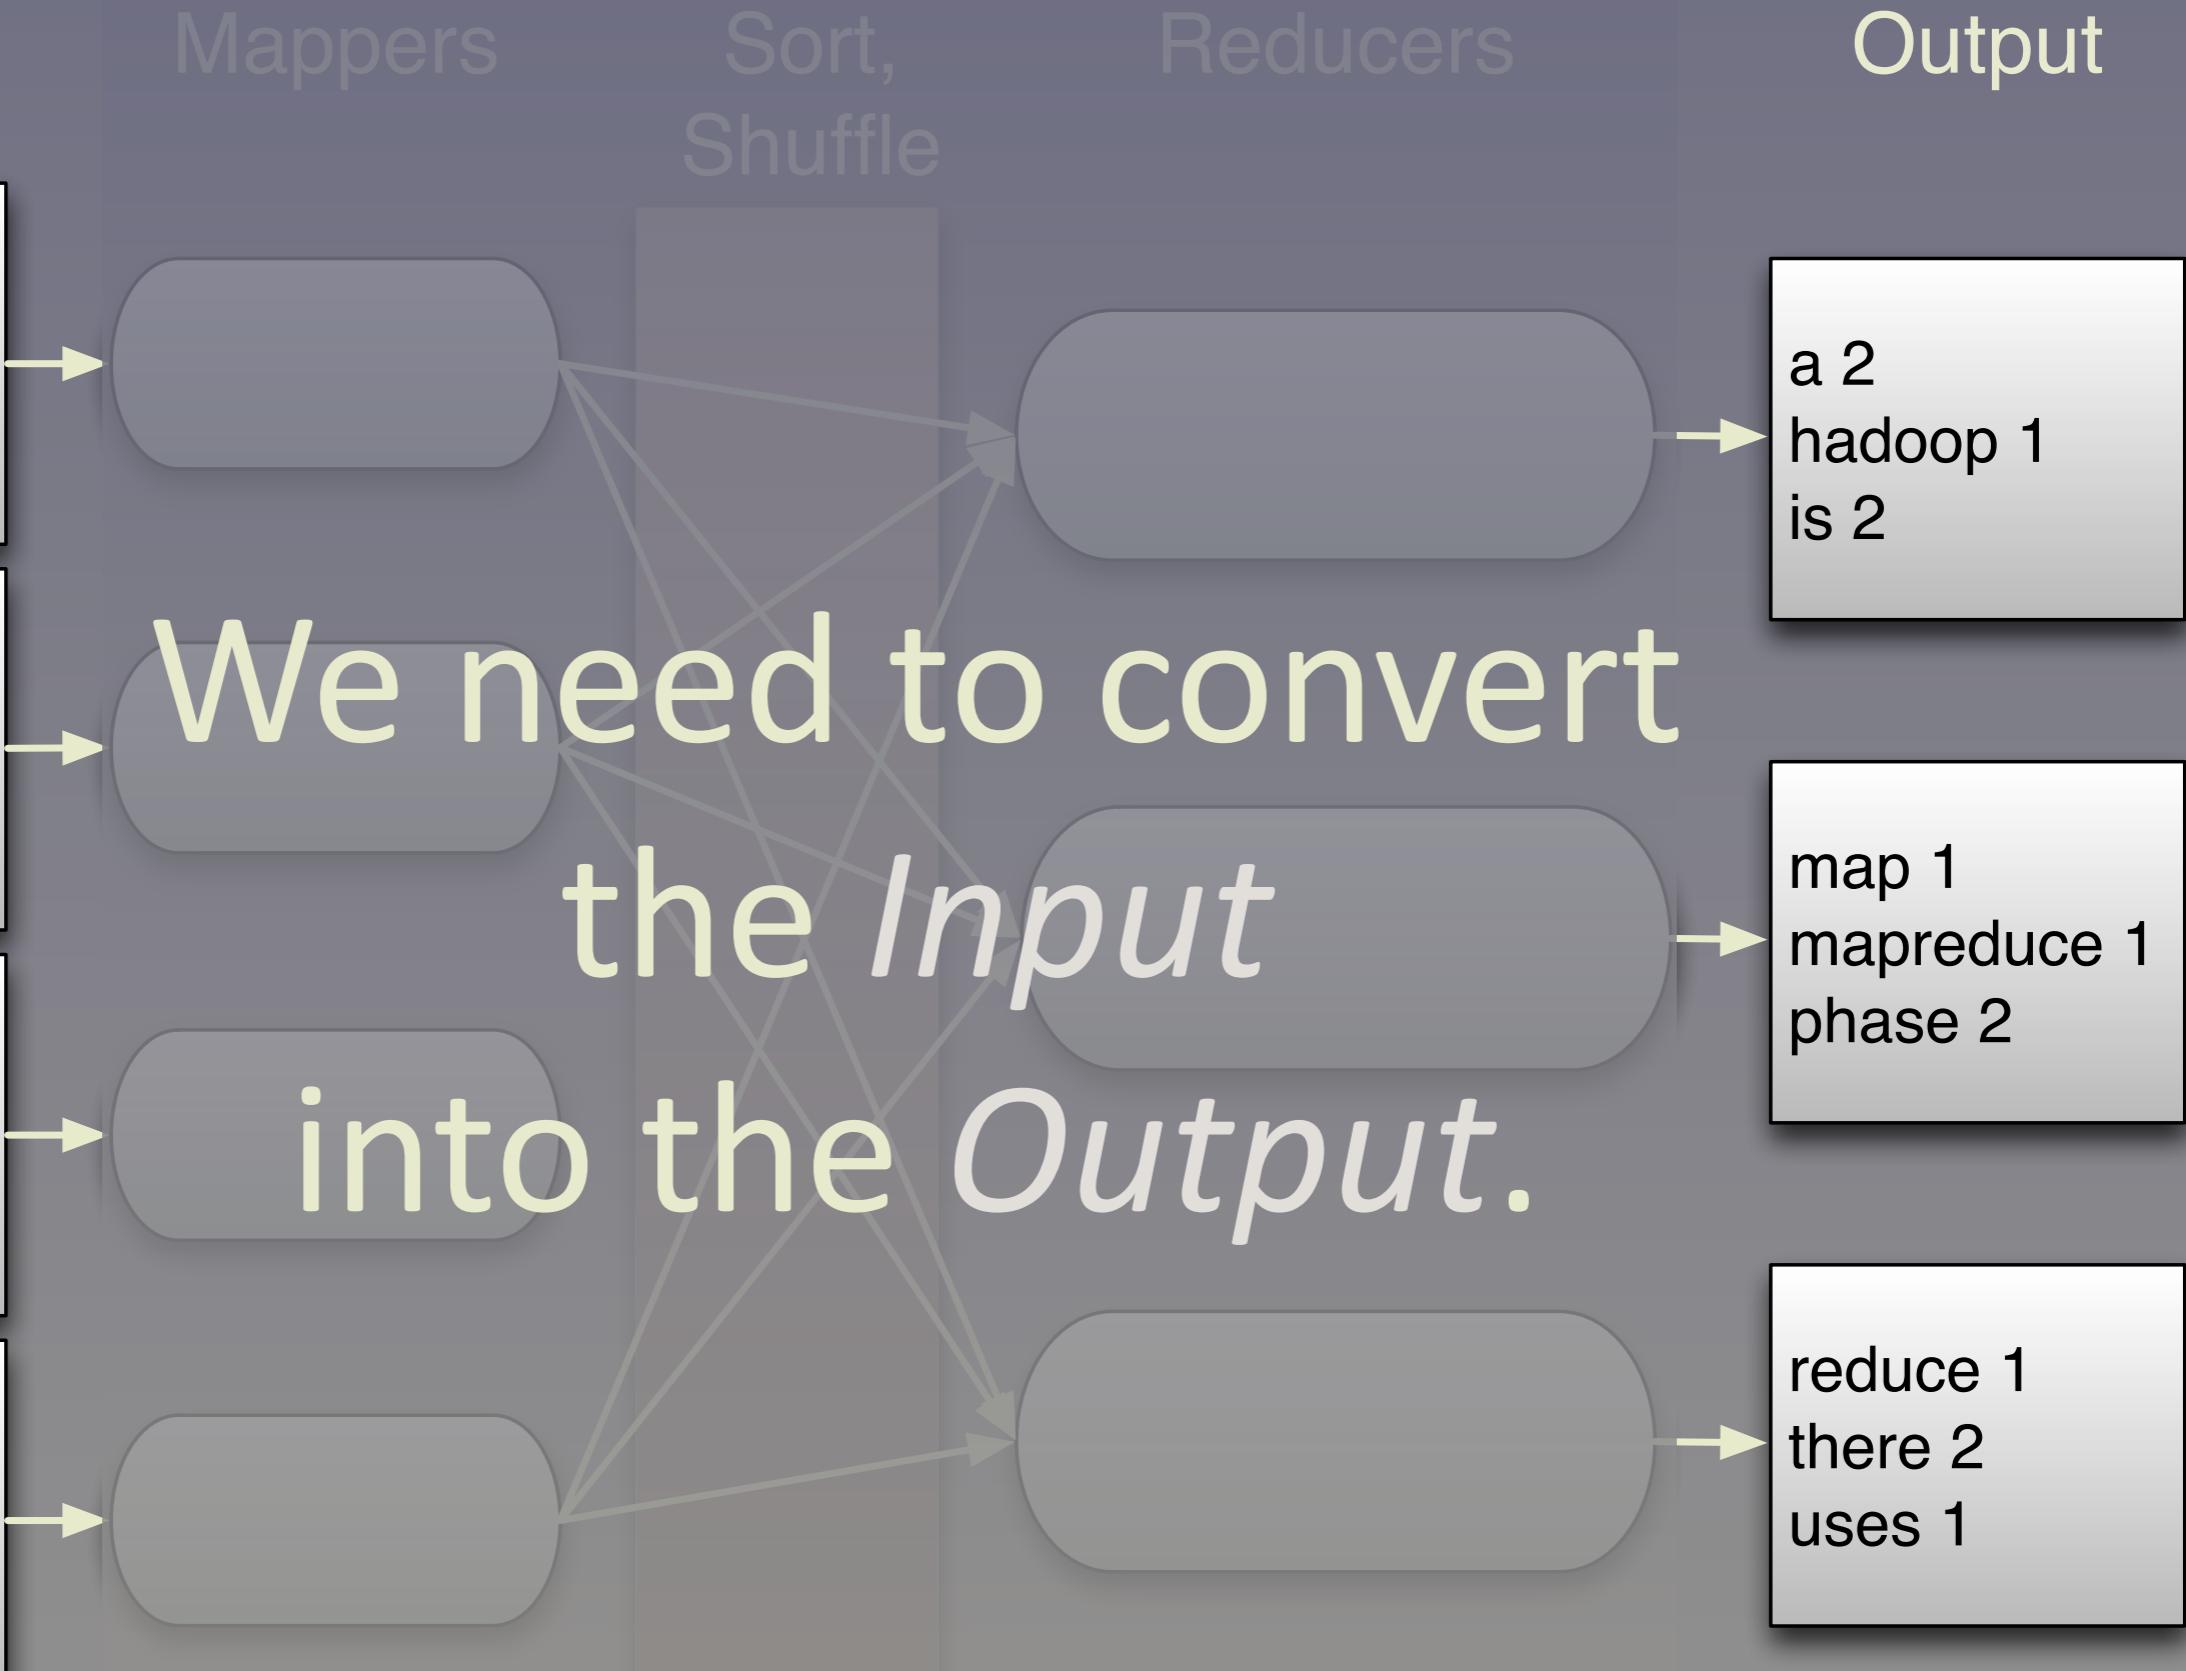
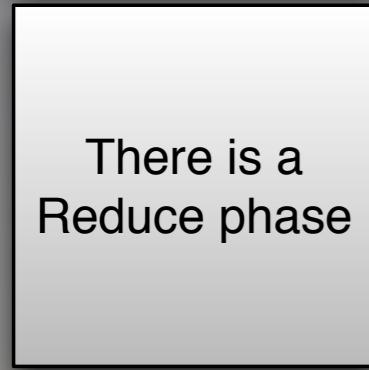
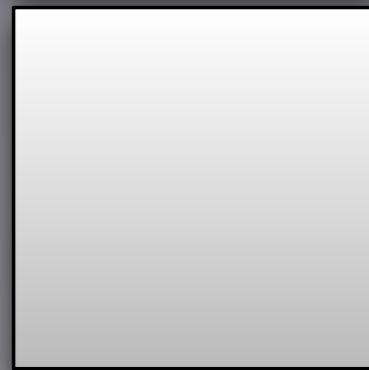
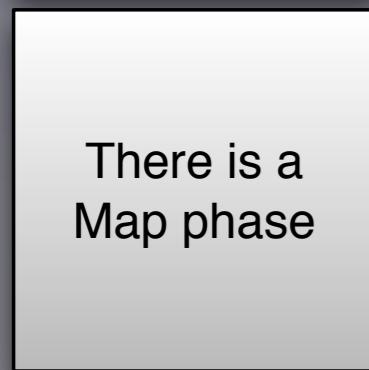
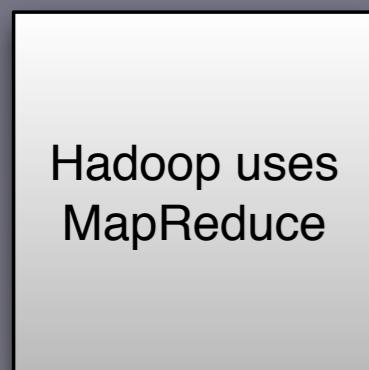
Input

Mappers

Sort,
Shuffle

Reducers

Output



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Four input documents, one left empty, the others with small phrases (for simplicity...). The word count output is on the right (we'll see why there are three output "documents"). We need to get from the input on the left-hand side to the output on the right-hand side.

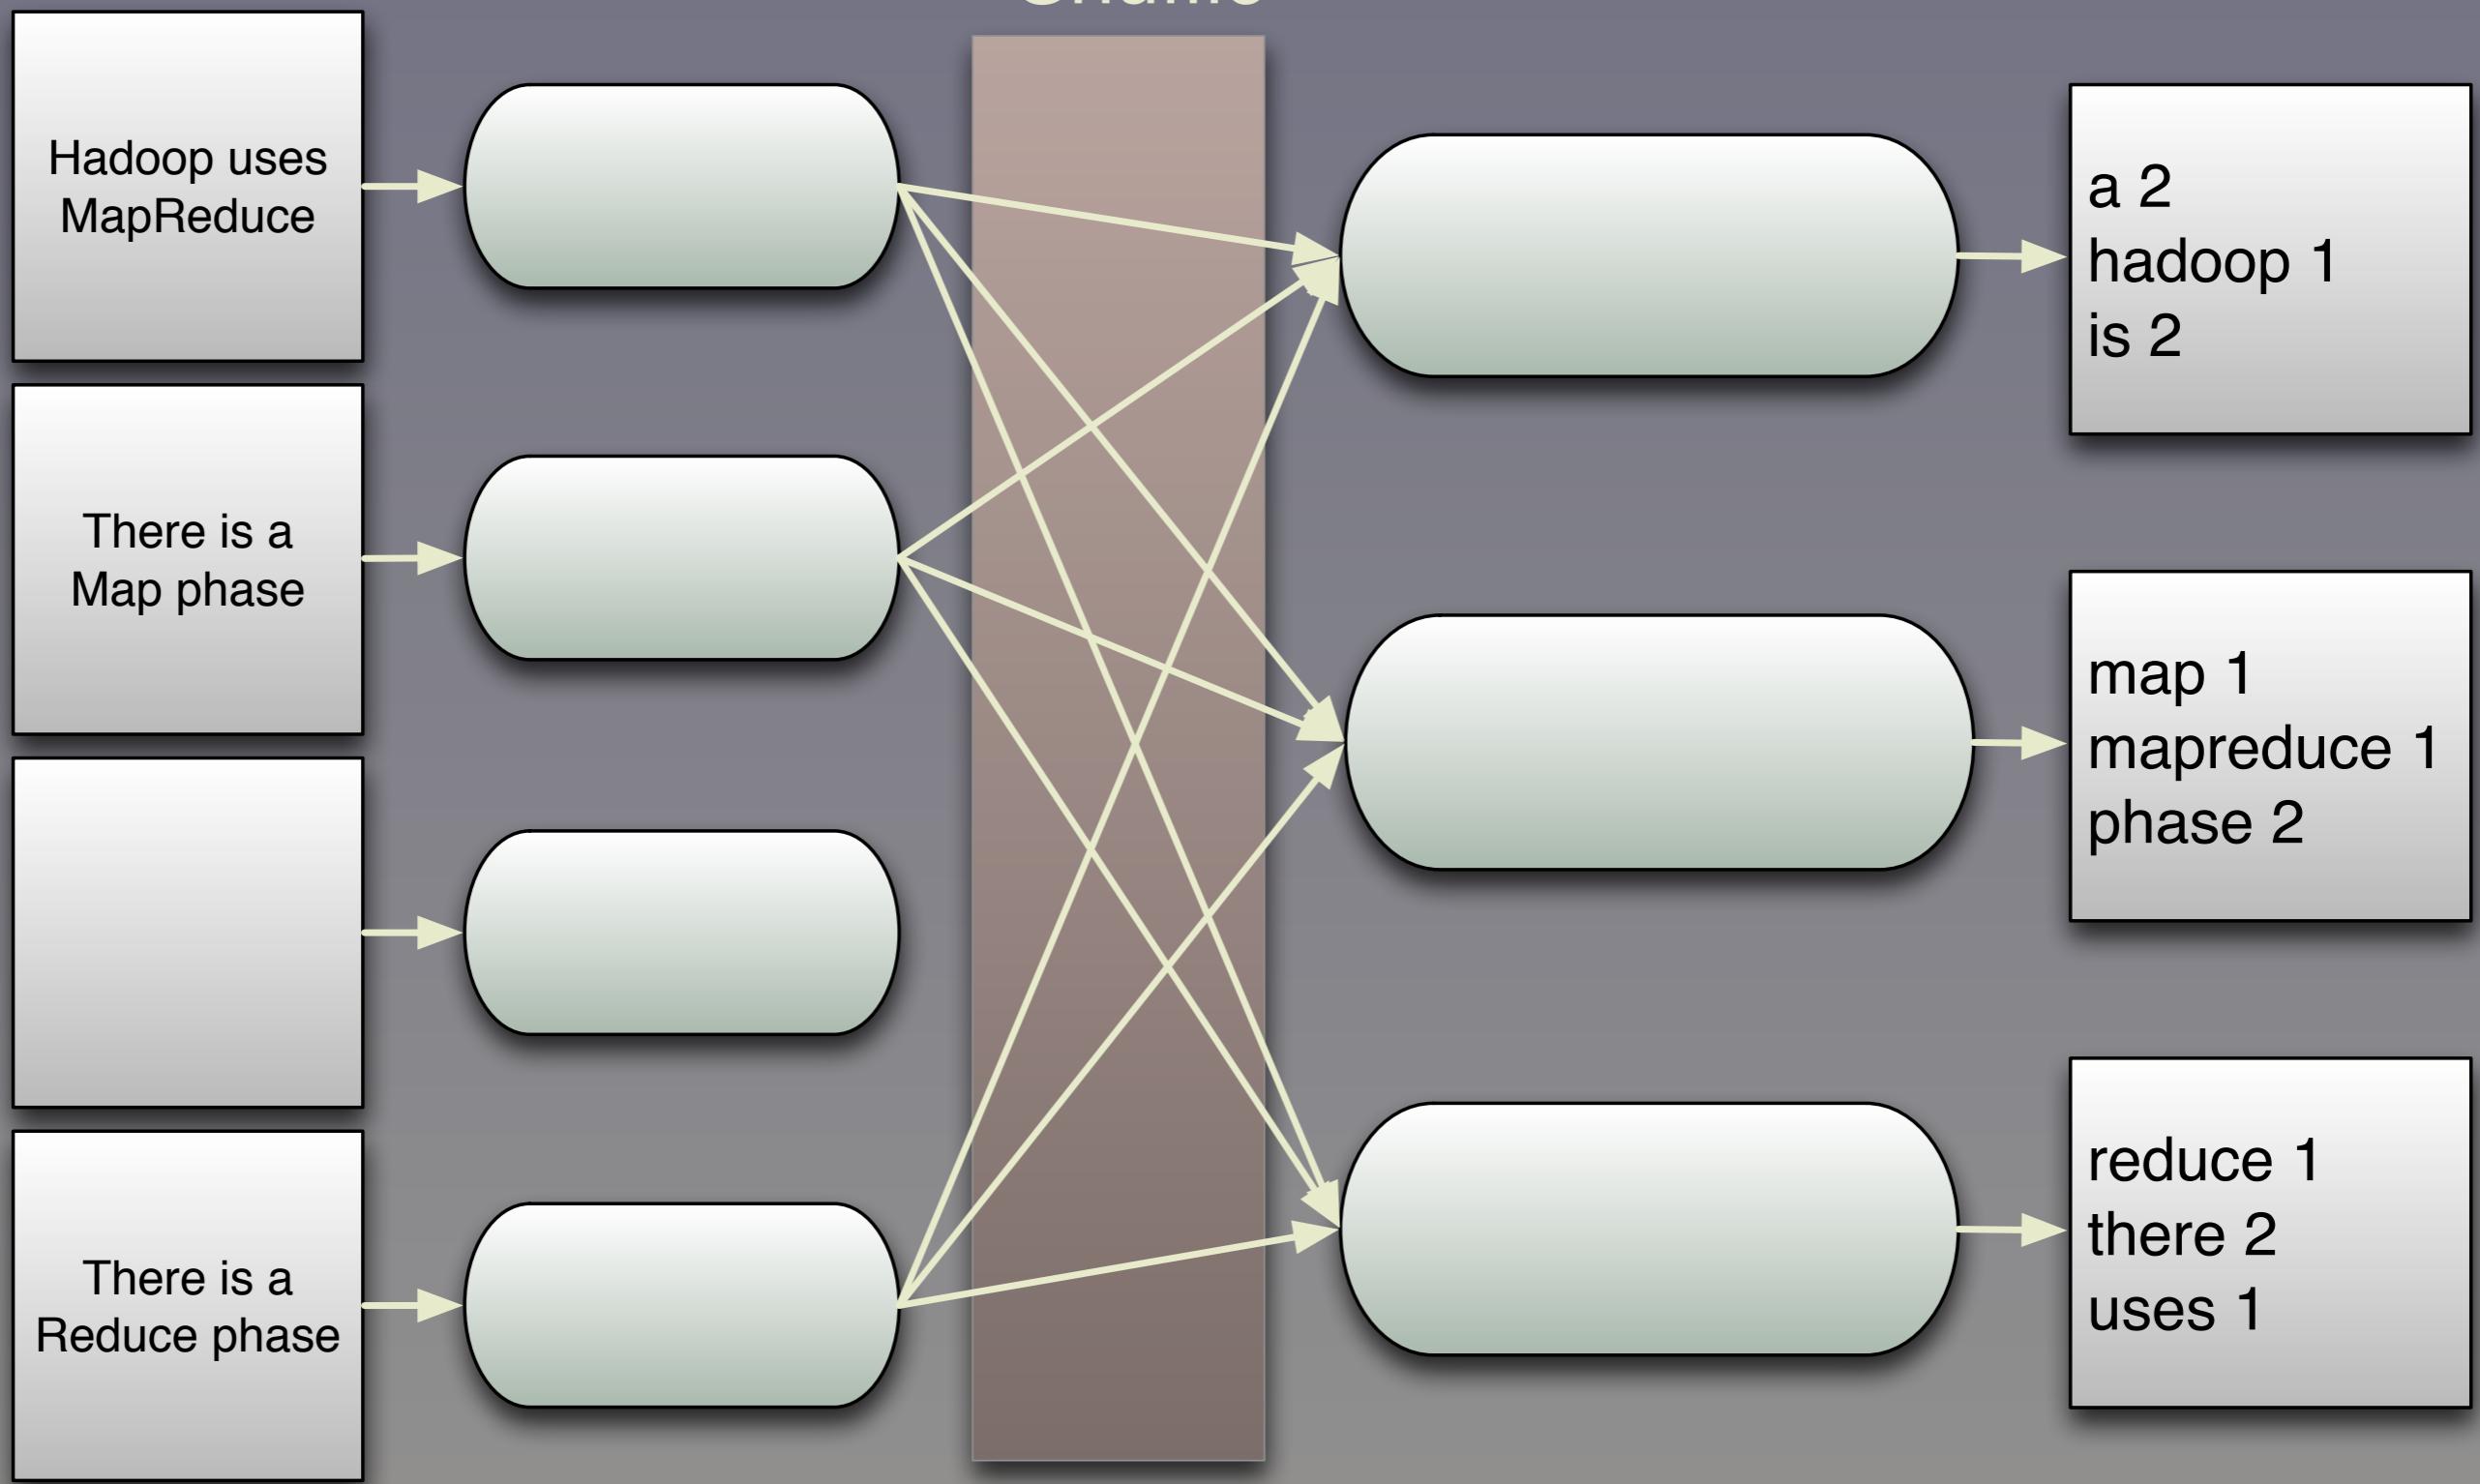
Input

Mappers

Sort, Shuffle

Reducers

Output

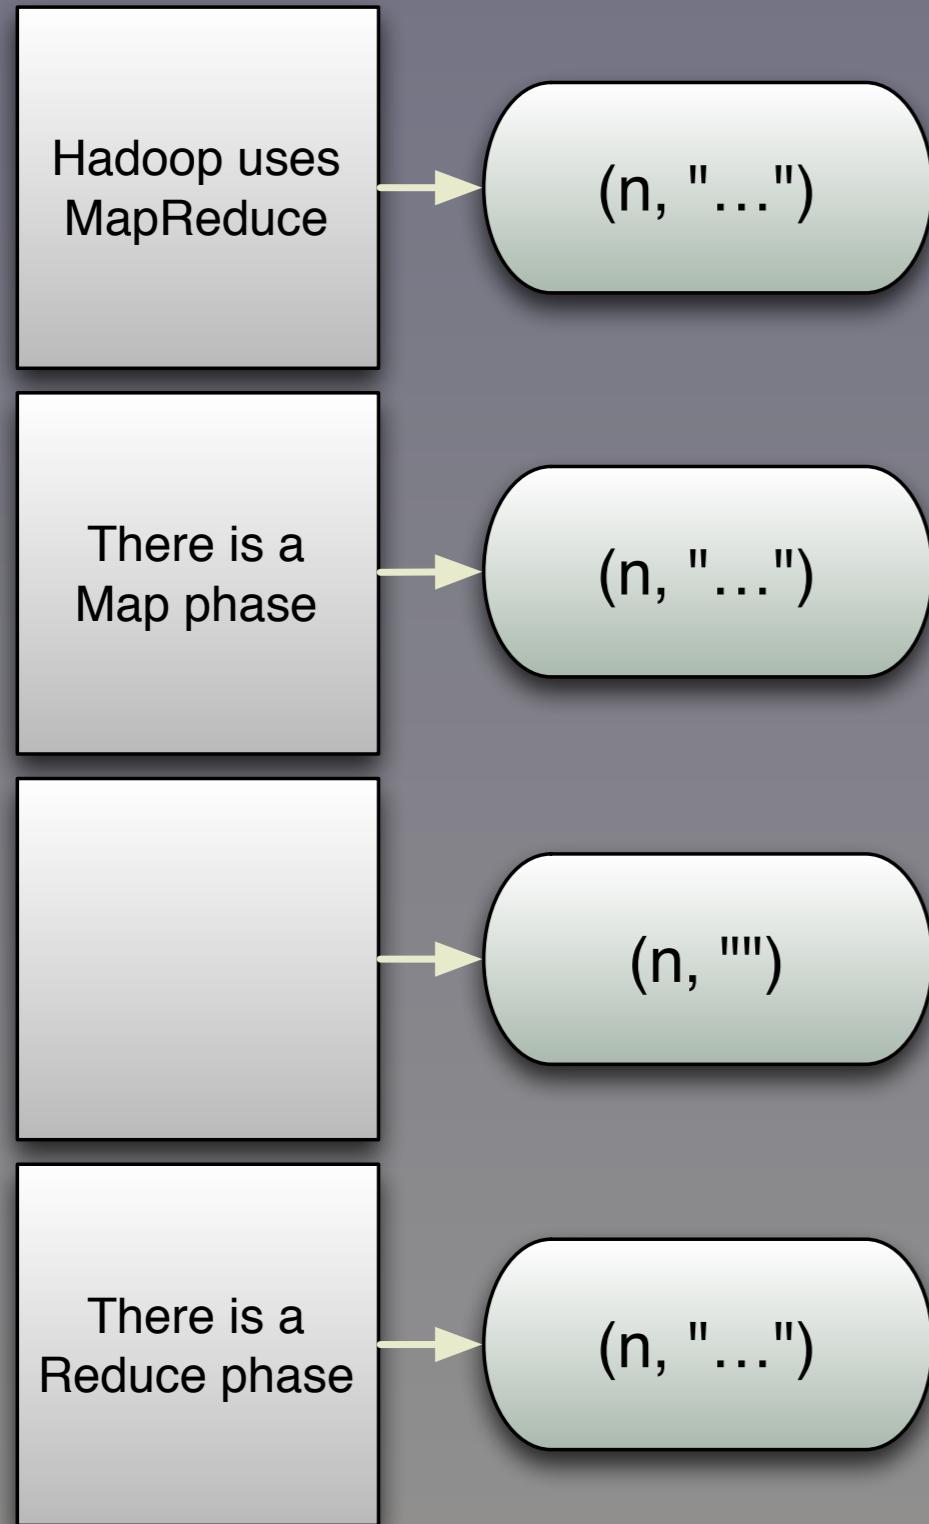


Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Here is a schematic view of the steps in Hadoop MapReduce. Each Input file is read by a single Mapper process (default: can be many-to-many, as we'll see later). The Mappers emit key-value pairs that will be sorted, then partitioned and "shuffled" to the reducers, where each Reducer will get all instances of a given key (for 1 or more values). Each Reducer generates the final key-value pairs and writes them to one or more files (based on the size of the output).

Input Mappers

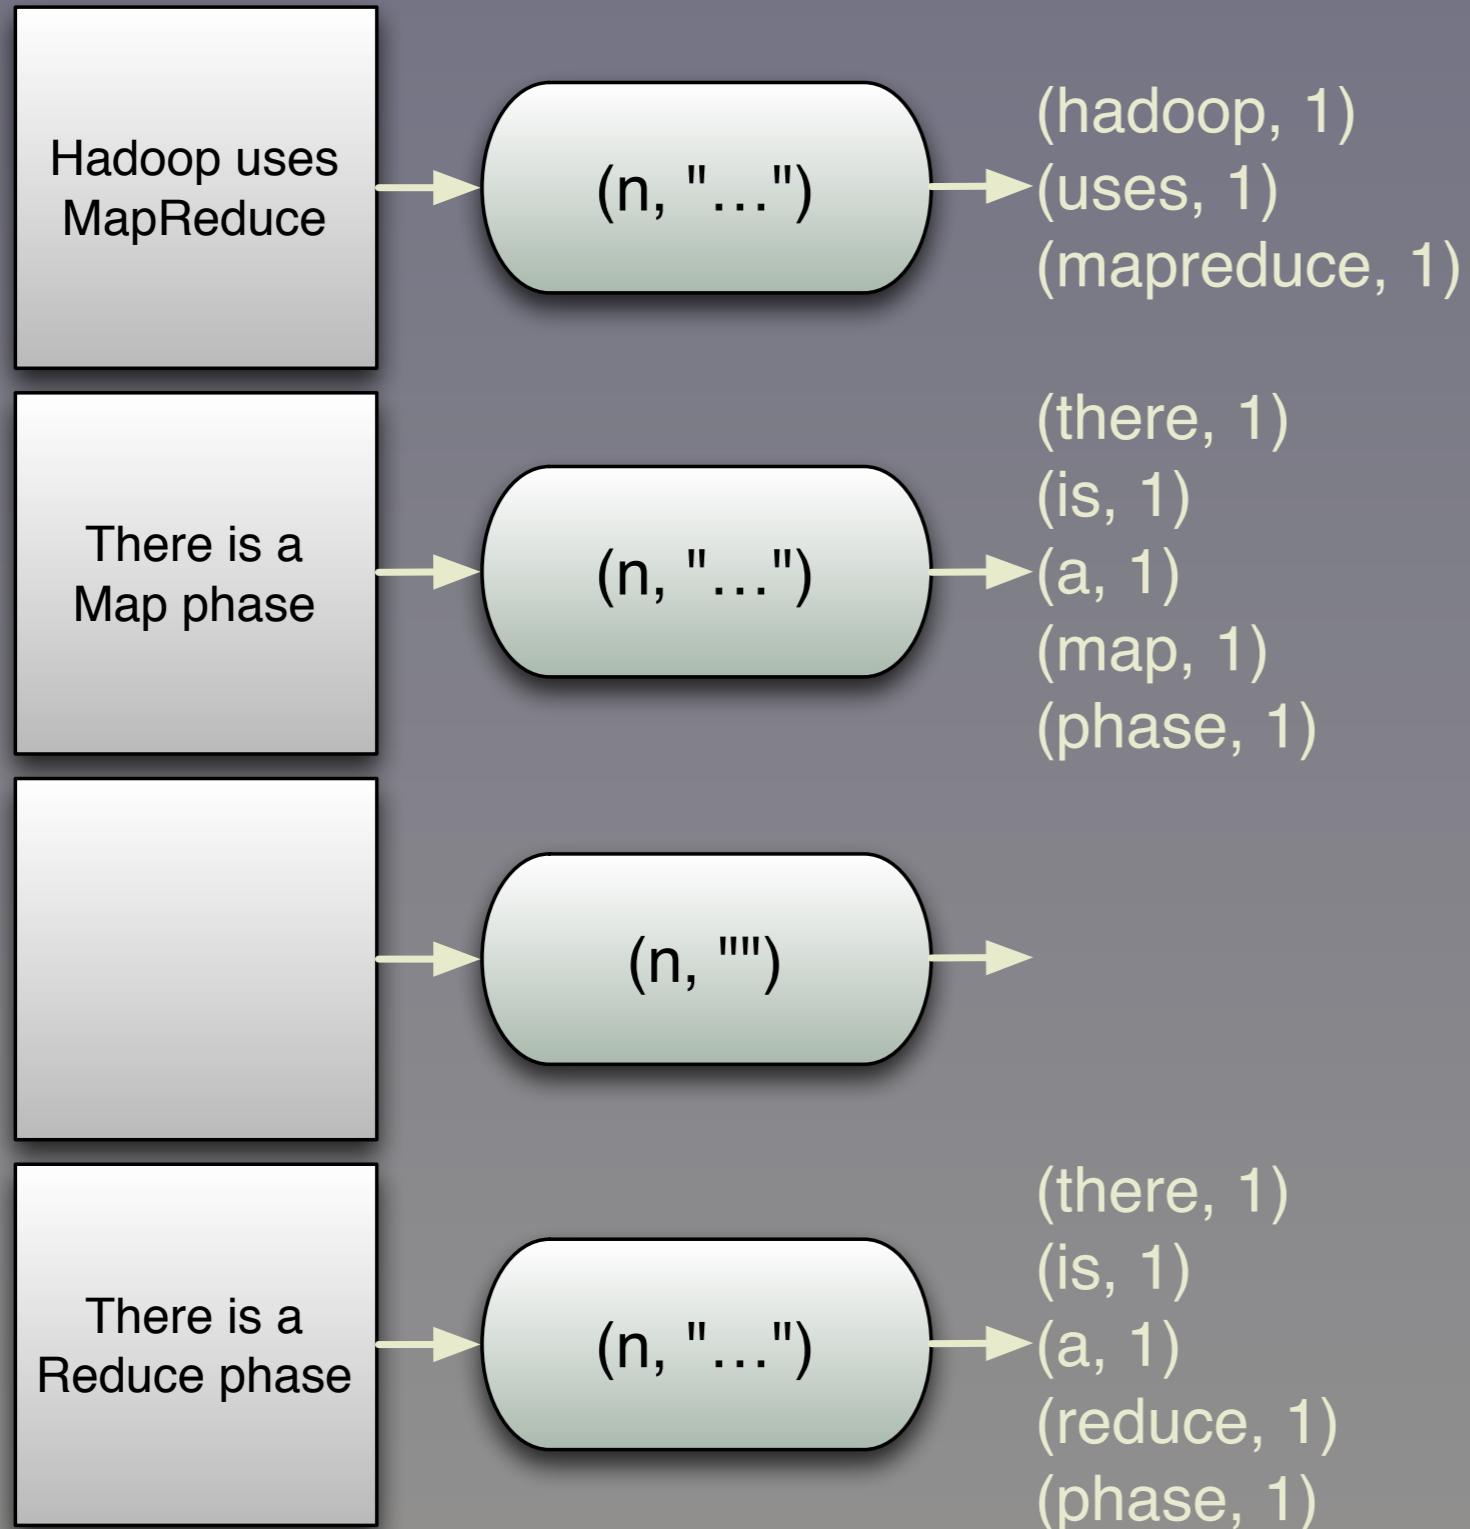


Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Each document gets a mapper. All data is organized into key-value pairs; each line will be a value and the offset position into the file will be the key, which we don't care about. I'm showing the document contents in the boxes and 1 mapper task (JVM process) per document. Large documents might get split to several mappers.
Each mapper will tokenize each line, one at a time, convert all words to lower case and count them...

Input Mappers



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

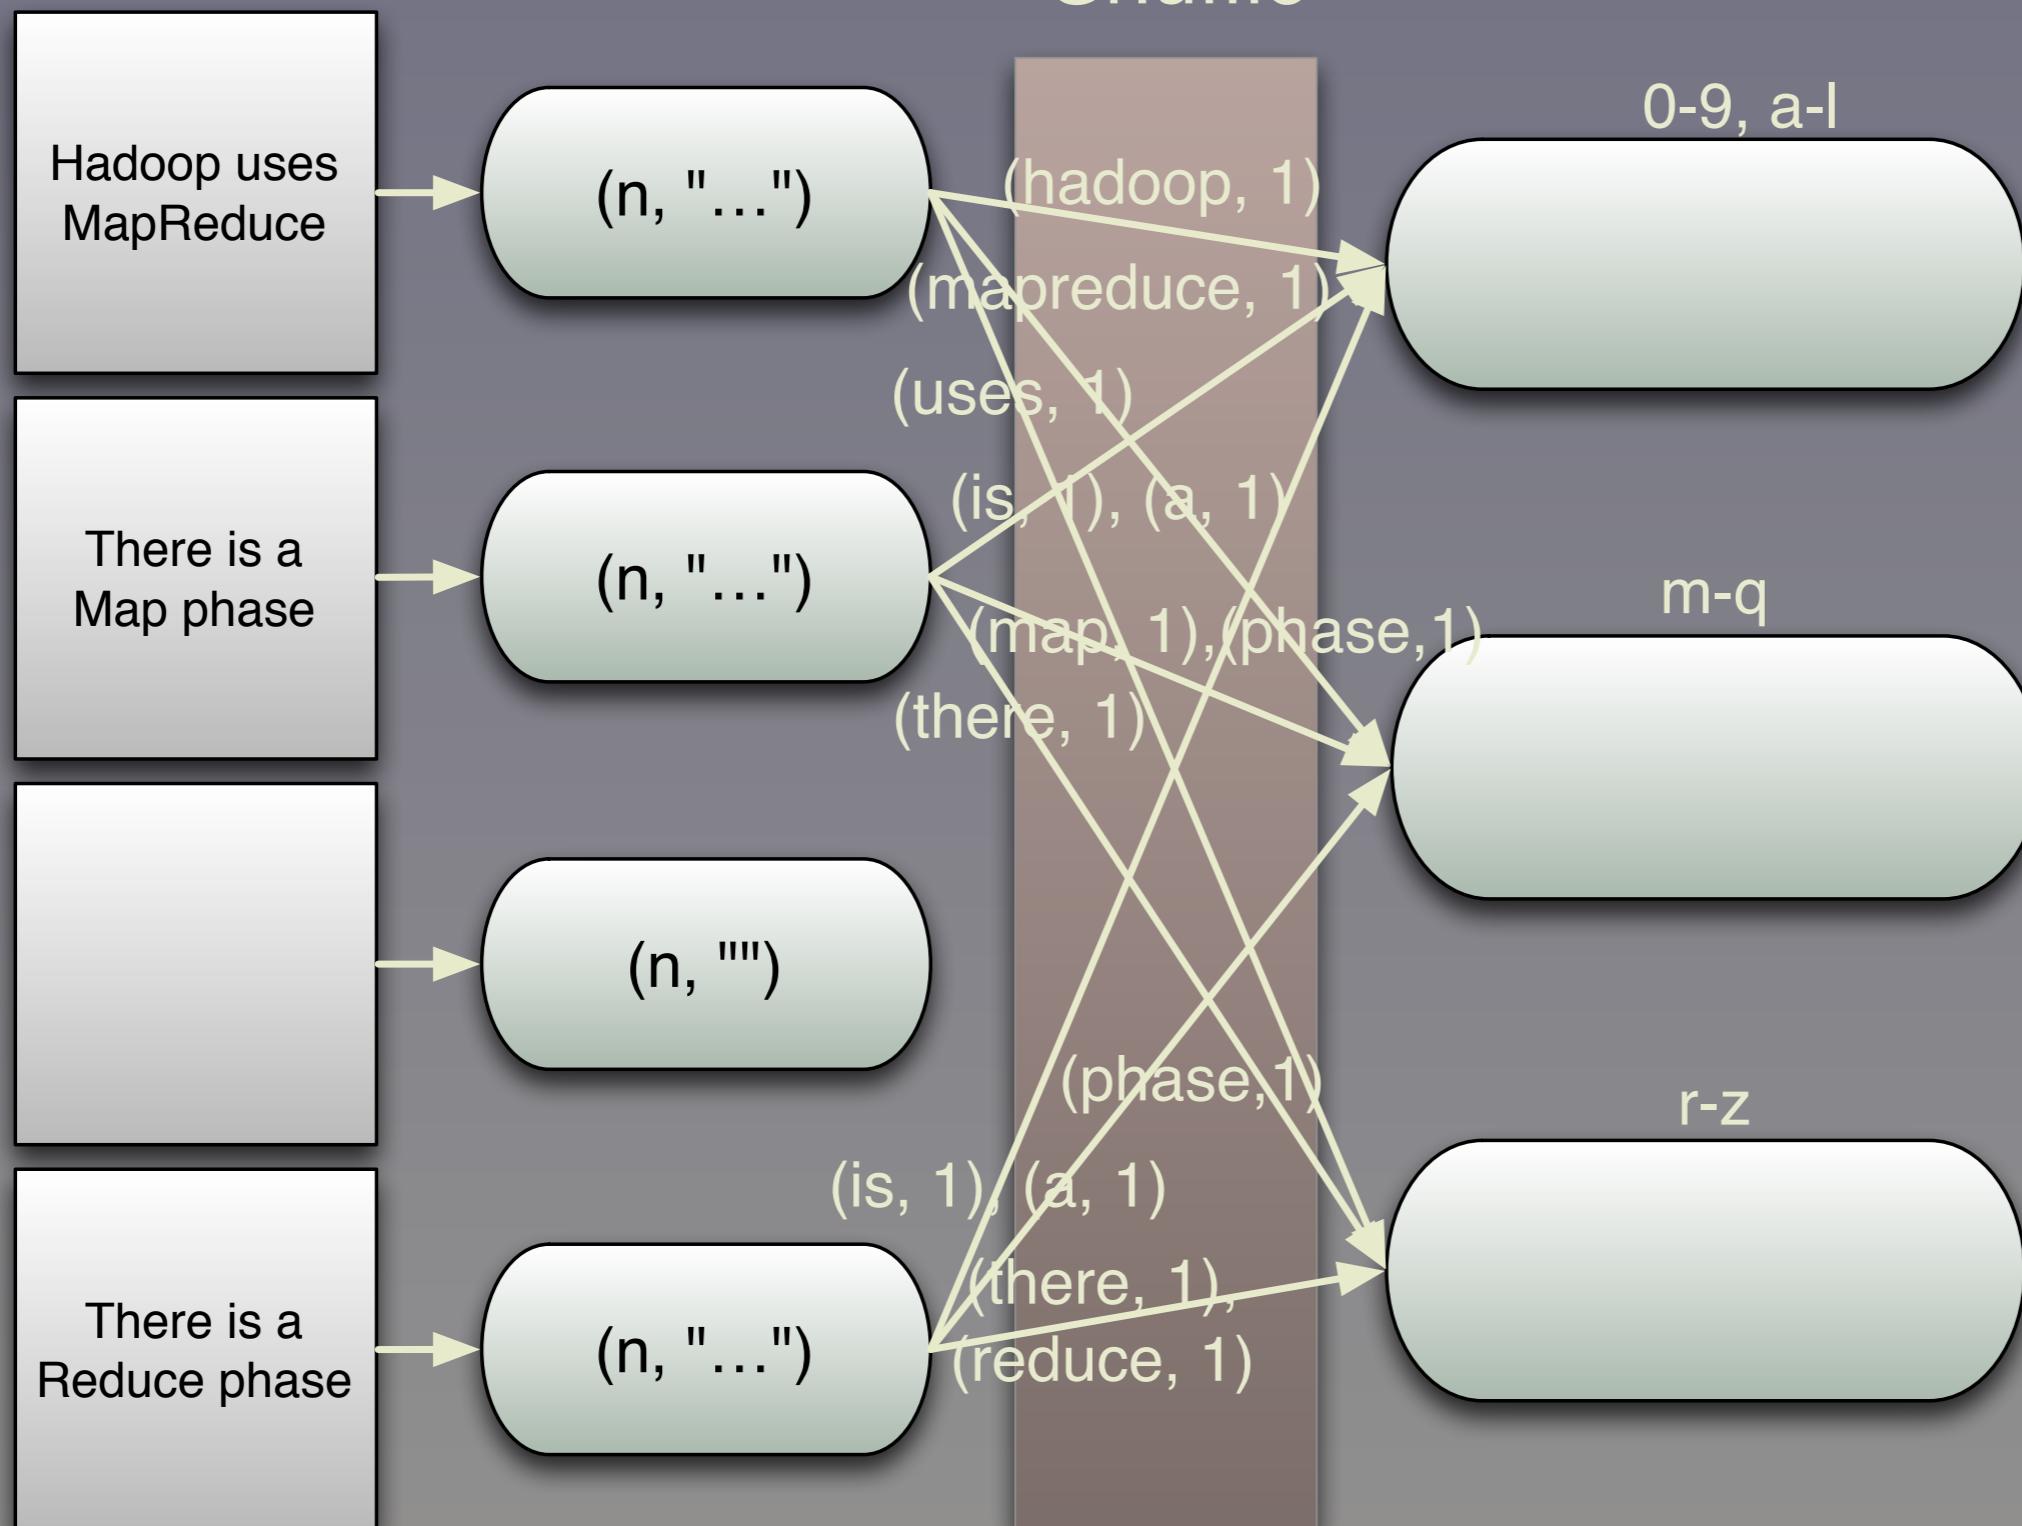
The mappers emit key-value pairs, where each key is one of the words, and the value is the count. In the most naive (but also most memory efficient) implementation, each mapper simply emits (word, 1) each time “word” is seen. However, this is IO inefficient! Note that the mapper for the empty doc. emits no pairs, as you would expect.

Input

Mappers

Sort, Shuffle

Reducers



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

The mappers themselves don't decide to which reducer each pair should be sent. Rather, the job setup configures what to do and the Hadoop runtime enforces it during the Sort/Shuffle phase, where the key-value pairs in each mapper are sorted by key (that is locally, not globally) and then the pairs are routed to the correct reducer, on the current machine or other machines.

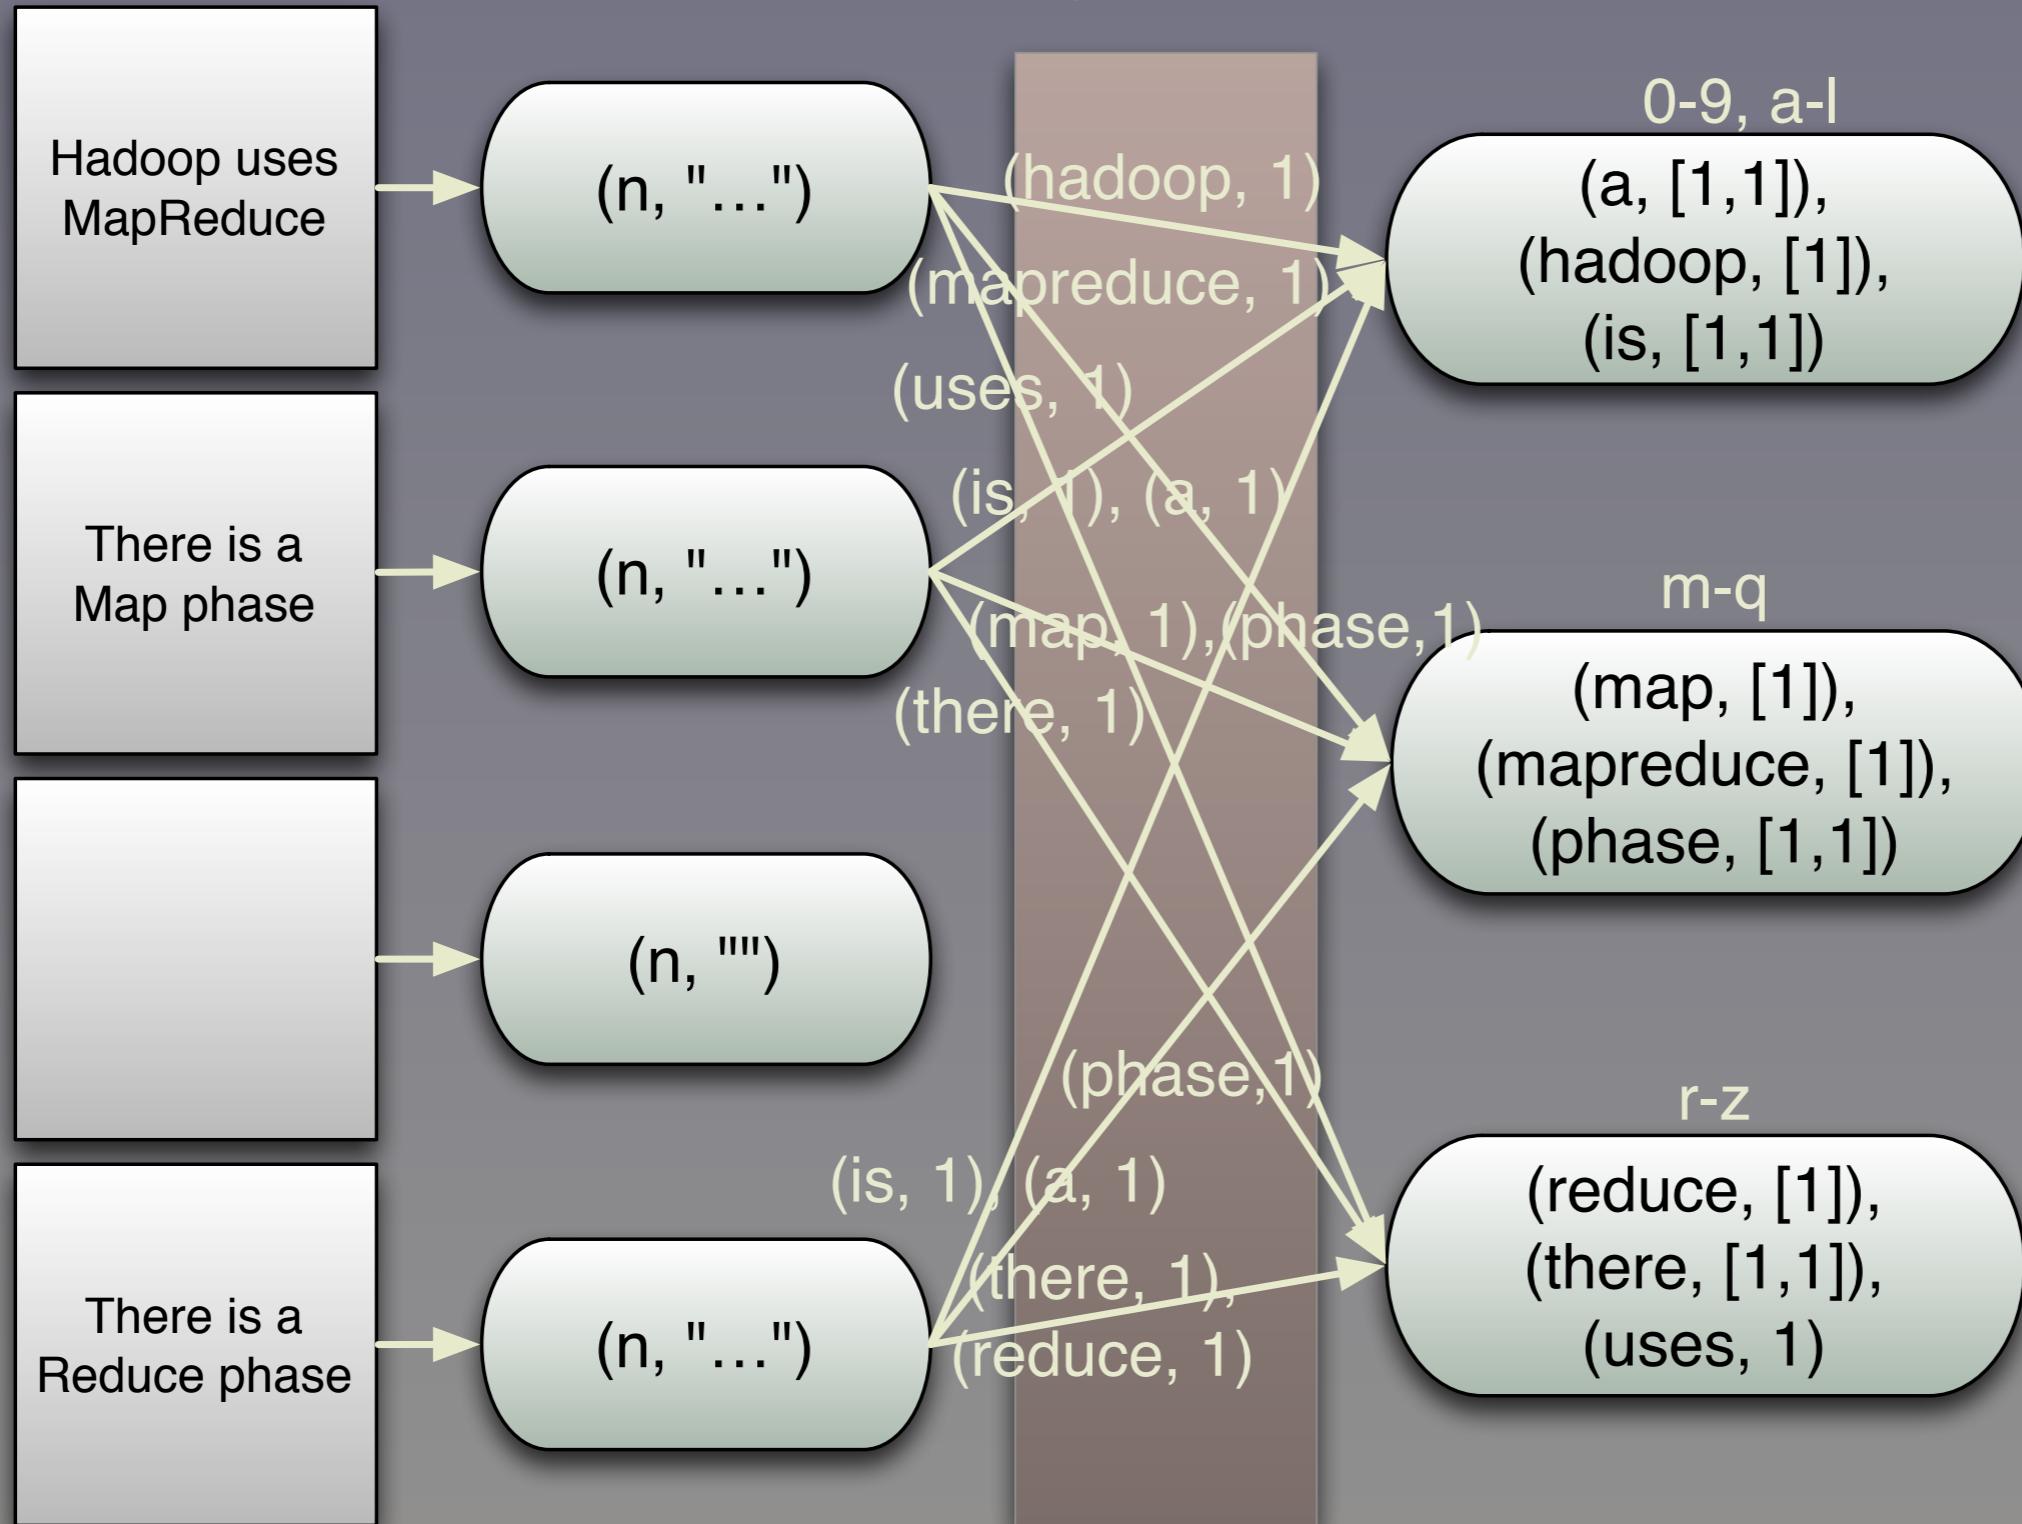
Note how we partitioned the reducers, by first letter of the keys. (By default, MR just hashes the keys and distributes them modulo # of reducers.)

Input

Mappers

Sort, Shuffle

Reducers



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Our reducers are passed each key and a collection of all the values for that key. (The MR framework creates this collection for us.)

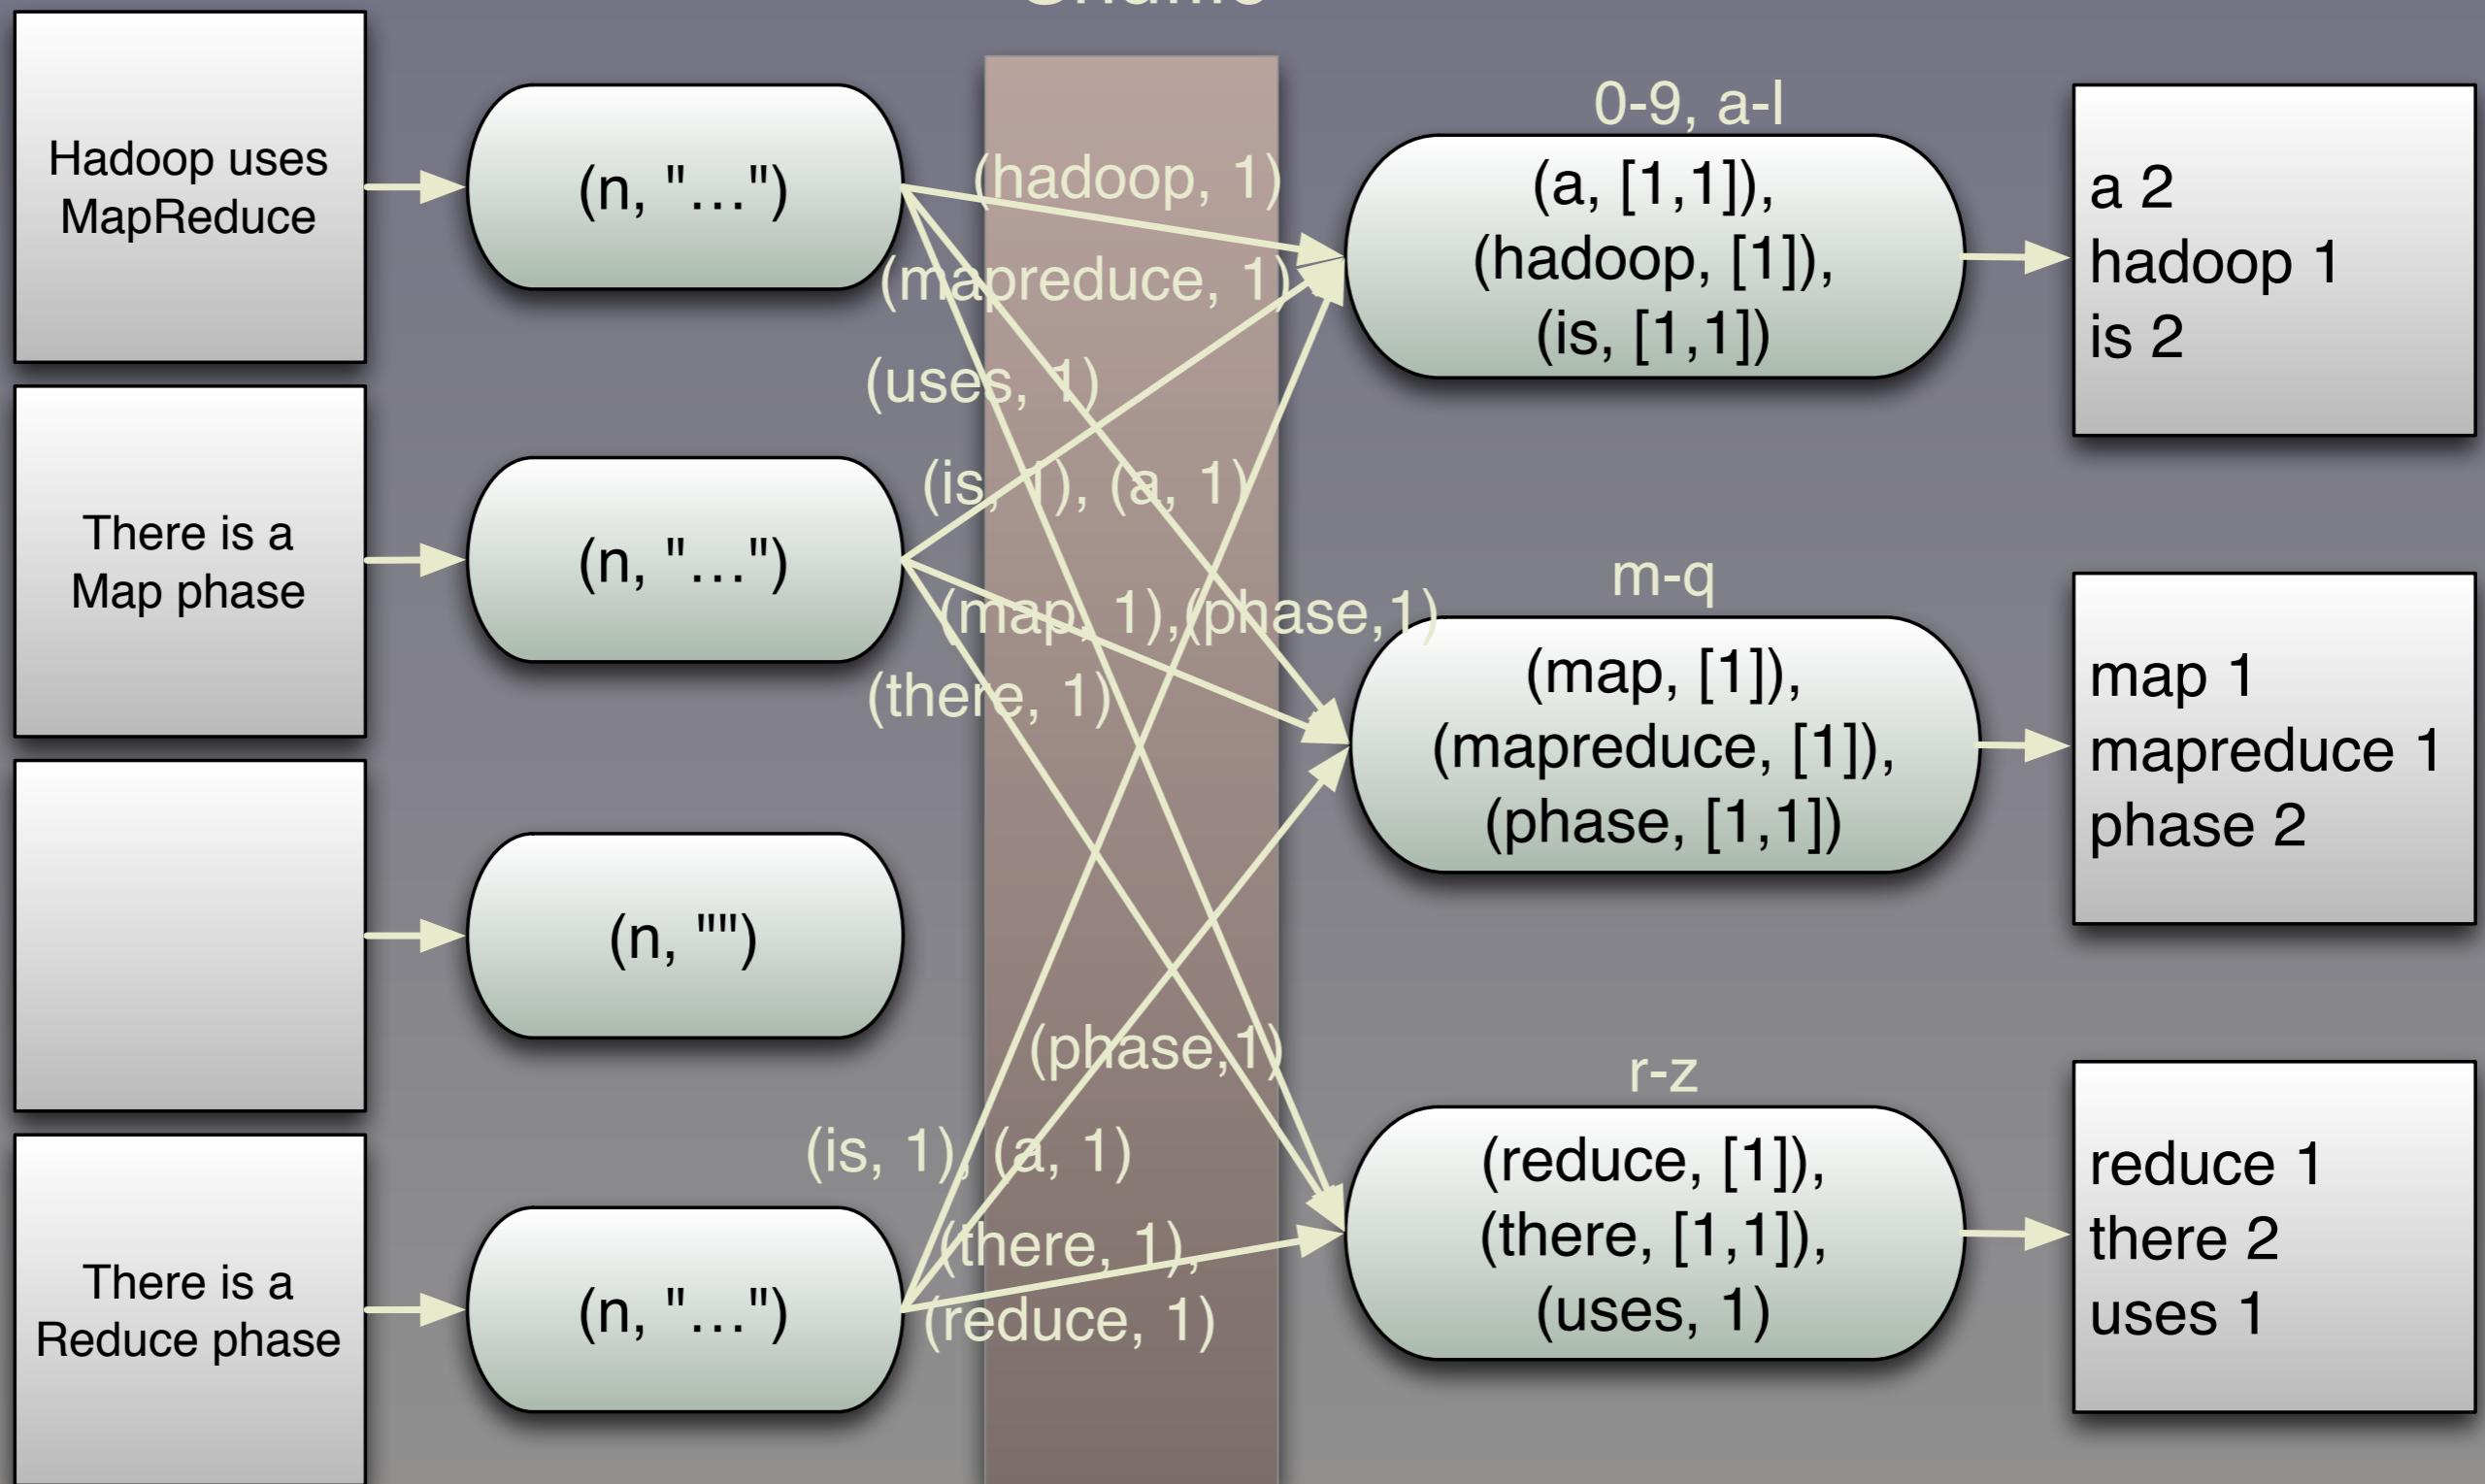
Input

Mappers

Sort, Shuffle

Reducers

Output



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

The final view of the WordCount process flow. The reducer just sums the counts and writes the output. The output files contain one line for each key (the word) and value (the count), assuming we're using text output. The choice of delimiter between key and value is up to you, but tab is common. (We'll discuss options as we go.)

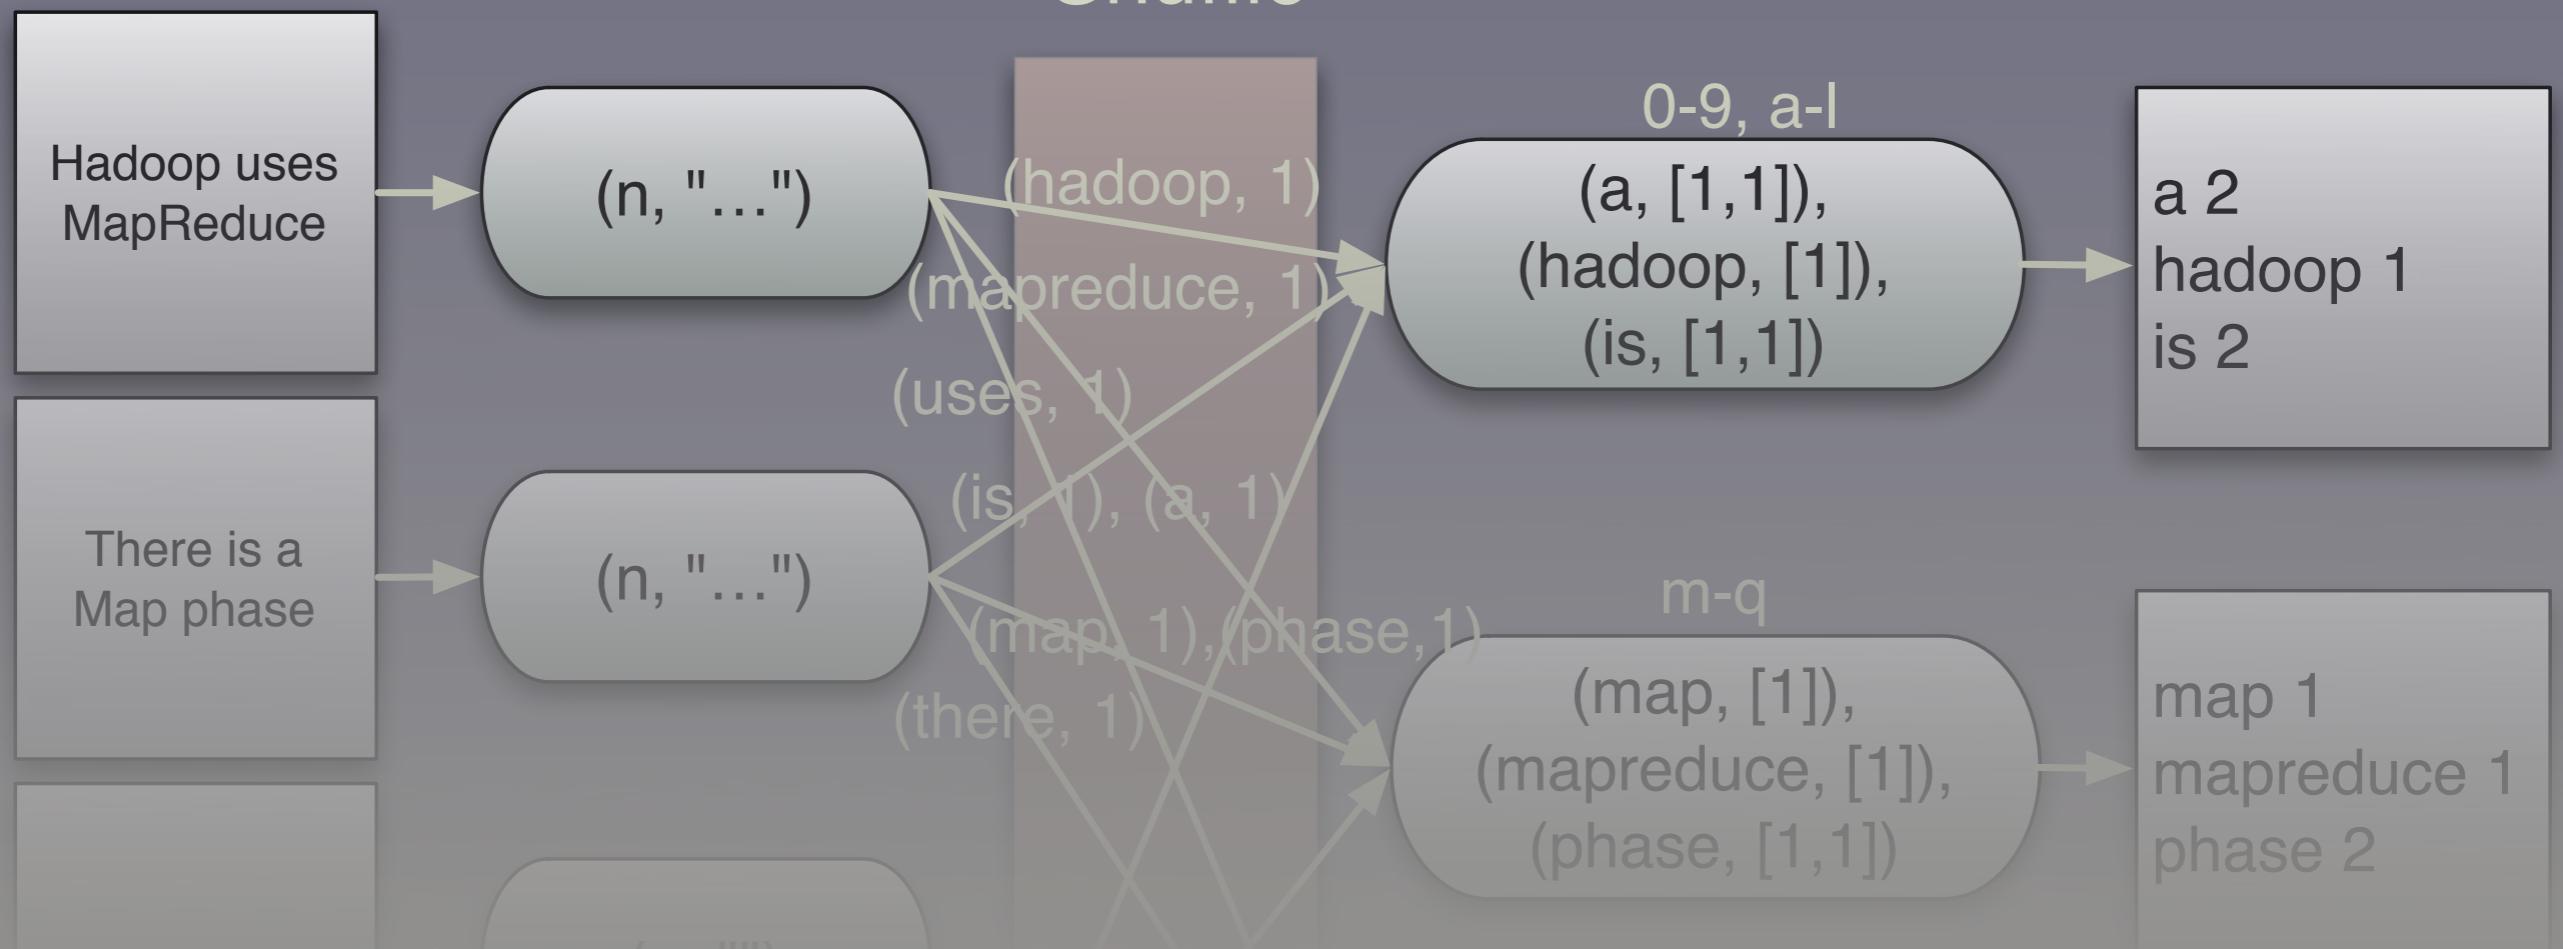
Input

Mappers

Sort, Shuffle

Reducers

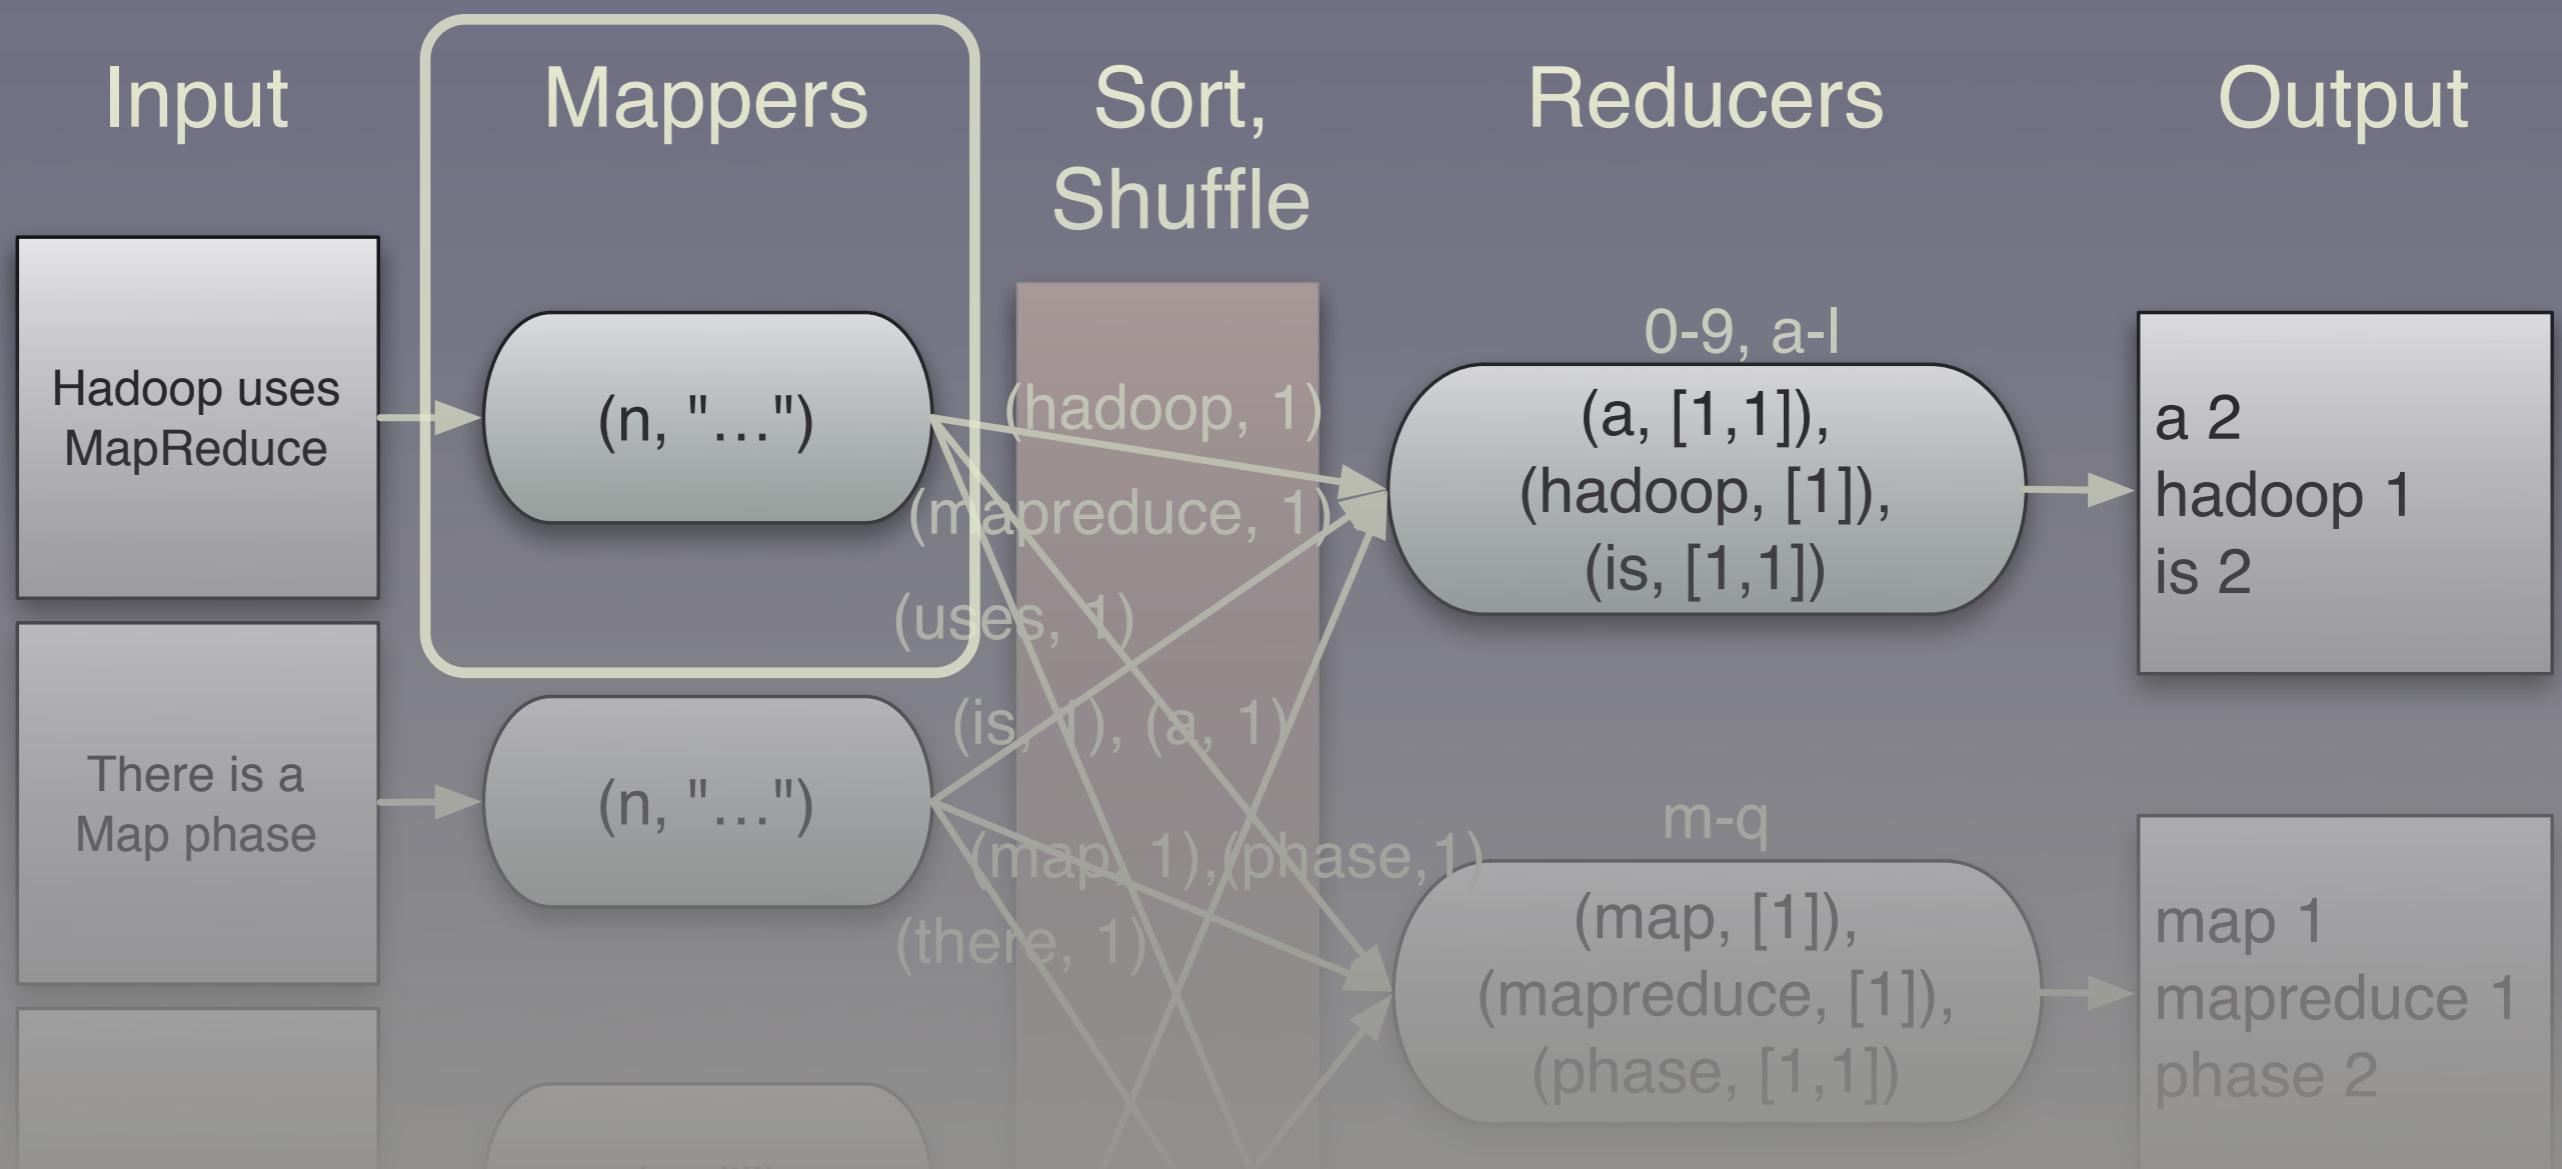
Output



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

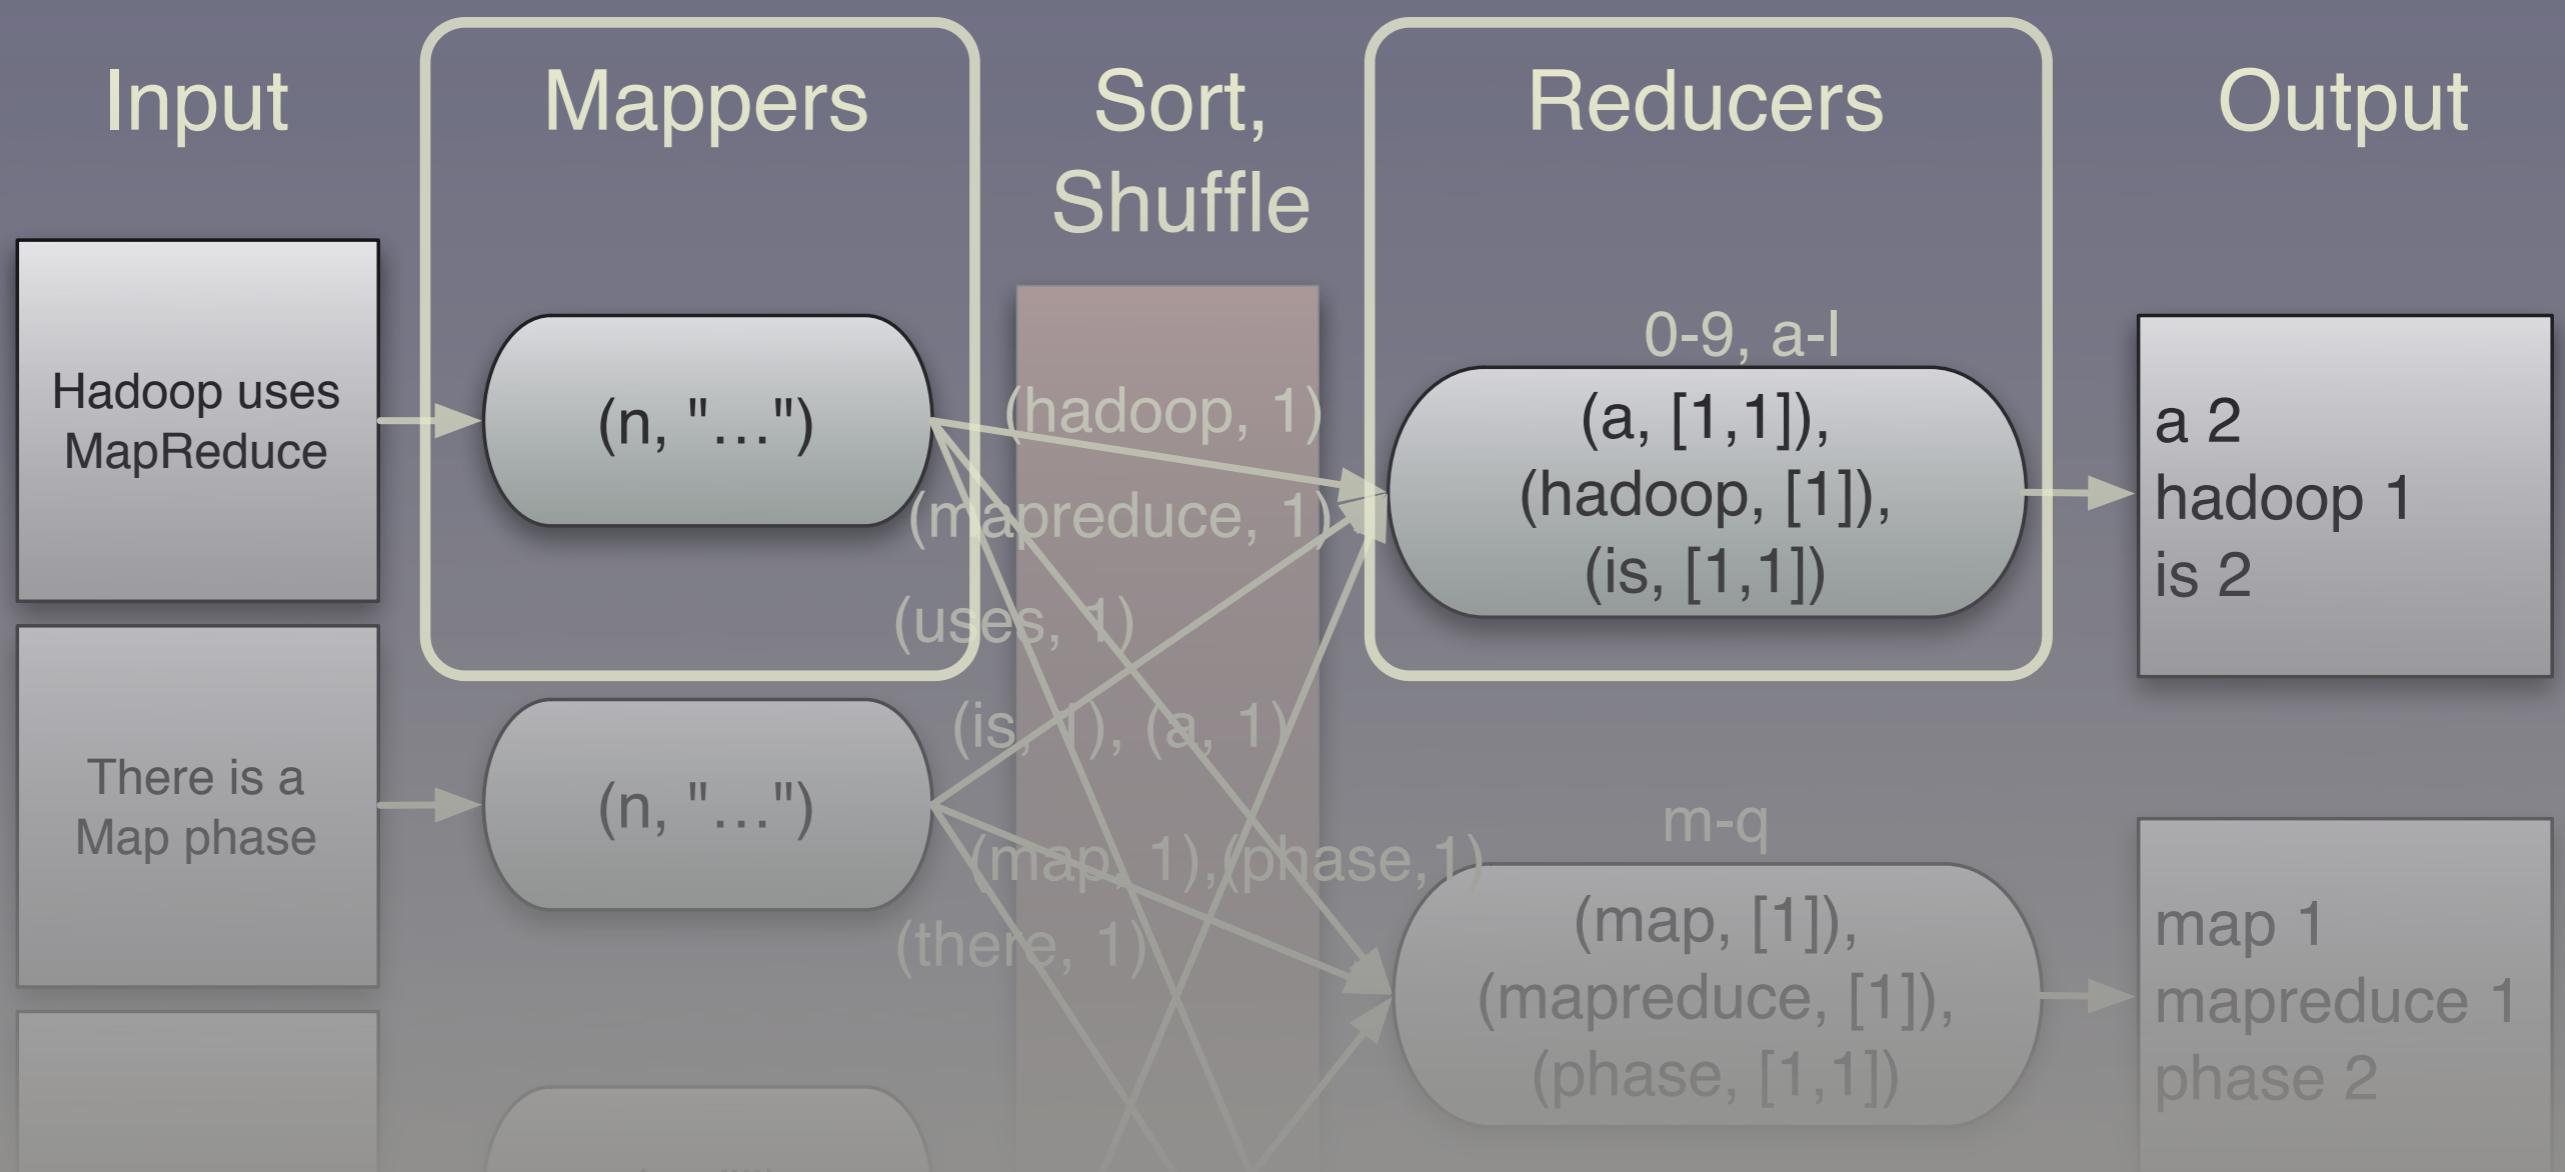
To recap, a “map” transforms one input to one output, but this is generalized in MapReduce to be one to 0-N. The output key-value pairs are distributed to reducers. The “reduce” collects together multiple inputs with the same key into



Map:

- Transform *one* input to *0-N* outputs.

To recap, a “map” transforms one input to one output, but this is generalized in MapReduce to be one to 0-N. The output key-value pairs are distributed to reducers. The “reduce” collects together multiple inputs with the same key into



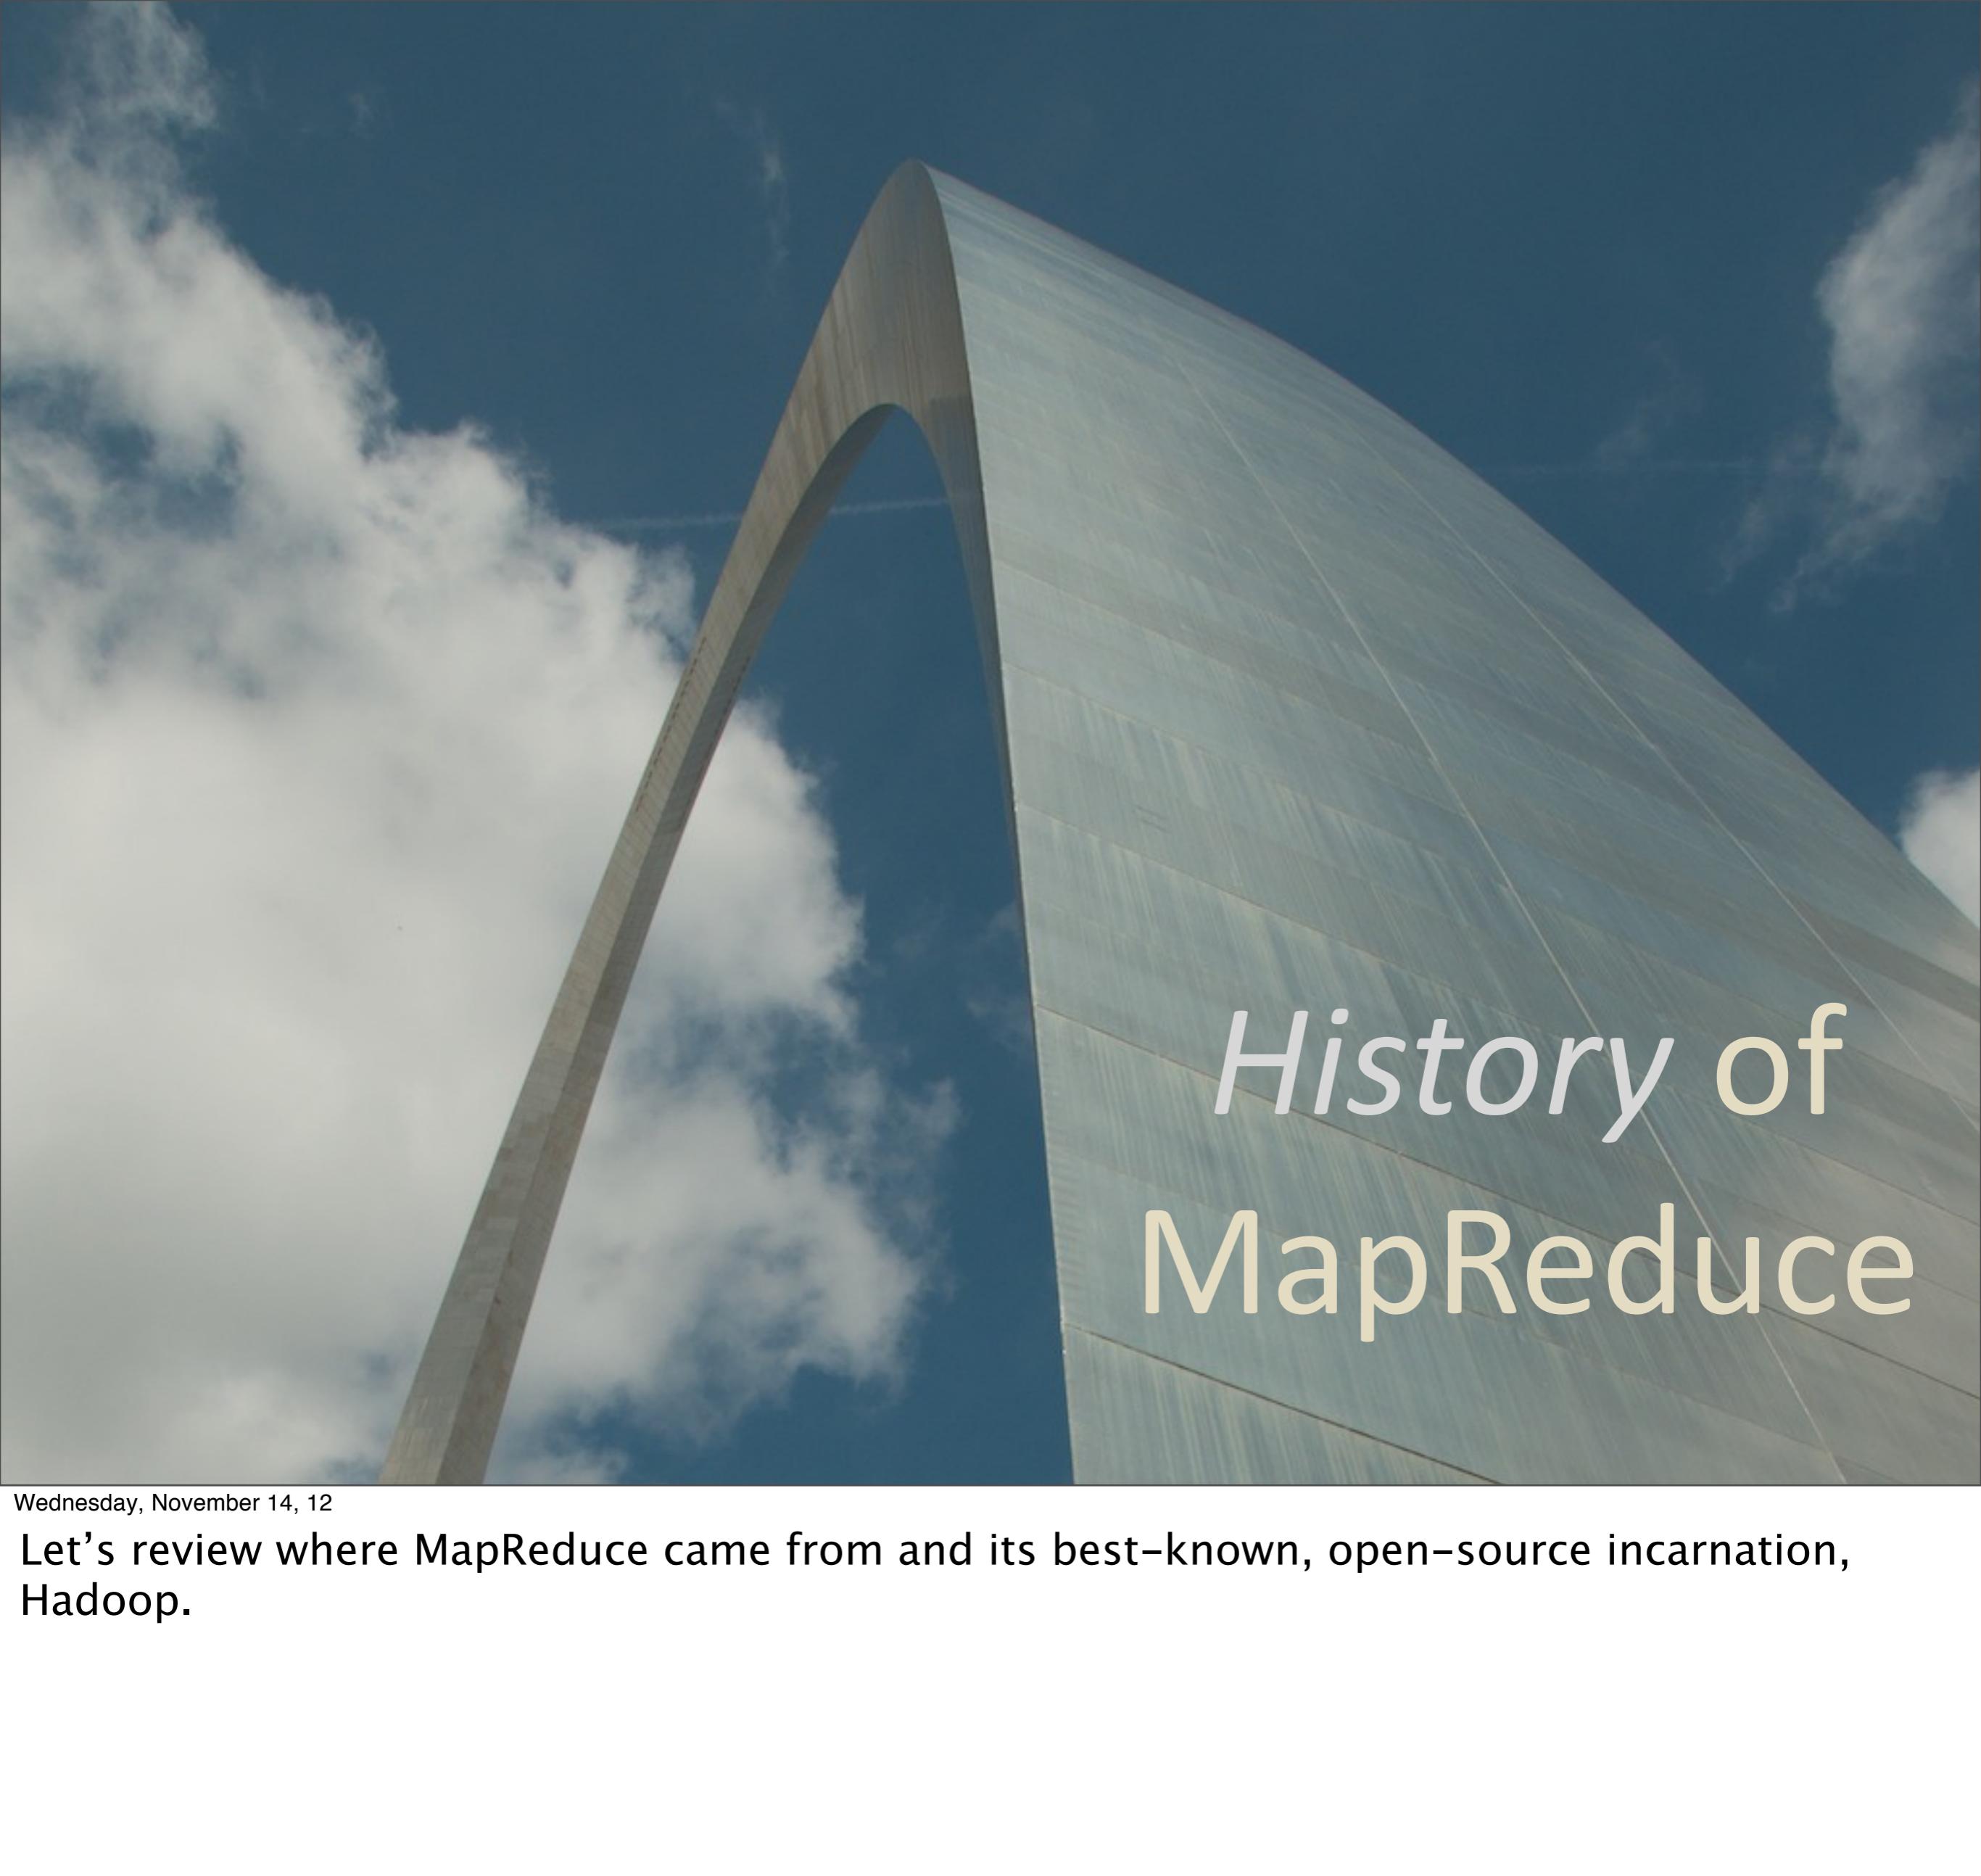
Map:

- Transform *one* input to *0-N* outputs.

Reduce:

- Collect *multiple* inputs into one output.

To recap, a “map” transforms one input to one output, but this is generalized in MapReduce to be one to 0-N. The output key-value pairs are distributed to reducers. The “reduce” collects together multiple inputs with the same key into

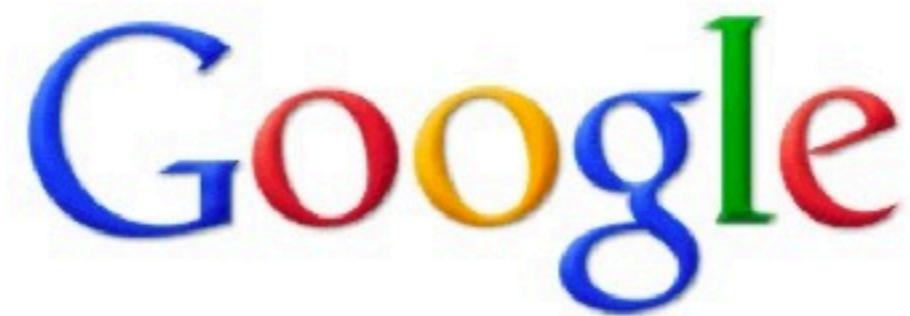
A photograph of a modern architectural structure, possibly a library or concert hall, featuring a large glass-enclosed entrance and a curved, light-colored facade. The sky is blue with scattered white clouds.

History of MapReduce

Wednesday, November 14, 12

Let's review where MapReduce came from and its best-known, open-source incarnation, Hadoop.

How would you *index the web?*

A screenshot of a Google search interface. The search bar at the top contains the partial query "What is the meanin". Below the search bar, a list of suggested queries is displayed, all starting with "what is the meaning of":

- what is the meaning of life
- what is the meaning of life 42
- what is the meaning of pumped up kicks
- what is the meaning of halloween
- what is the meaning of love
- what is the meaning of labor day
- what is the meaning of slope
- what is the meaning of homecoming
- what is the meaning of my last name
- what is the meaning of a promise ring

At the bottom of the interface are two buttons: "Google Search" and "I'm Feeling Lucky".

How would you *index the web?*

Google

what is the meaning of life

Search

About 48,900,000 results (0.26 seconds)

Everything

Images

Maps

Videos

News

Shopping

[Meaning of life - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Meaning_of_life +1
The **meaning of life** constitutes a philosophical question concerning the purpose and significance of life or existence in general. This concept can be expressed ...
[Questions - Western philosophical perspectives - East Asian philosophy](#)

[The Meaning of Life](#)
users.aristotle.net/~diogenes/meaning1.htm +1
Why do you want to know the **meaning of life**? Often people ask this question when they really want the answer to some other question. Let's try and get those ...

Wednesday, November 14, 12

You ask a *phrase* and
the *search engine* finds
the *best match* in
billions of web pages.

Actually, Google
computes the *index*
that *maps* terms to
pages in *advance*.

Actually, Google
computes the *index*
that *maps* terms to
pages in *advance*.

Google's famous *Page Rank* algorithm.

Google File System

for Storage

2003

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung

Google*

ABSTRACT

We have designed and implemented the Google File System, a scalable distributed file system for large distributed data-intensive applications. It provides fault tolerance while running on inexpensive commodity hardware, and it delivers high aggregate performance to a large number of clients.

While sharing many of the same goals as previous distributed file systems, our design has been driven by observations of our application workloads and technological envi-

1. INTRODUCTION

We have designed and implemented the Google File System (GFS) to meet the rapidly growing demands of Google's data processing needs. GFS shares many of the same goals as previous distributed file systems such as performance, scalability, reliability, and availability. However, its design has been driven by key observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier

Wednesday, November 14, 12

A distributed file system provides horizontal scalability and resiliency when file blocks are duplicated around the cluster.

MapReduce for Computation

2004

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to par-

Wednesday, November 14, 12

A distributed file system provides horizontal scalability and resiliency when file blocks are duplicated around the cluster.

About this time, *Doug Cutting*, the creator of *Lucene* was working on *Nutch*...

Wednesday, November 14, 12

Lucene is an open-source text search engine. Nutch is an open source web crawler.

He implemented clean-
room versions of
MapReduce and *GFS*...

By 2006 , they became
part of a separate
Apache project,
called *Hadoop*.

Wednesday, November 14, 12

The name comes from a toy, stuffed elephant that Cutting's son owned at the time.

By 2006 , they became
part of a separate
Apache project,
called *Hadoop*.



Wednesday, November 14, 12

The name comes from a toy, stuffed elephant that Cutting's son owned at the time.

Benefits of MapReduce



Wednesday, November 14, 12

The best way to
approach *Big Data* is to
scale *Horizontally*.

Hadoop Design Goals

Wednesday, November 14, 12

Maximizing disk and network I/O is so important, because it's such a bottleneck, that it is a core design goal of Hadoop (both MR and HDFS). It affects the features and performance of everything in the stack, including high-level programming tools!

Hadoop Design Goals

Maximize I/O!!

Wednesday, November 14, 12

Maximizing disk and network I/O is so important, because it's such a bottleneck, that it is a core design goal of Hadoop (both MR and HDFS). It affects the features and performance of everything in the stack, including high-level programming tools!

Hadoop Design Goals

Maximize I/O!!

Oh, and run on *server-class, commodity hardware.*

Wednesday, November 14, 12

Maximizing disk and network I/O is so important, because it's such a bottleneck, that it is a core design goal of Hadoop (both MR and HDFS). It affects the features and performance of everything in the stack, including high-level programming tools!

As a result, Hadoop is
great for batch mode
data crunching.

MapReduce and its Discontents



Wednesday, November 14, 12

Is MapReduce the end of the story? Does it meet all our needs? Let's look at two problems...

#1

It's hard to *implement*
many *Algorithms*
in *MapReduce*.

Wednesday, November 14, 12

Even word count is not “obvious”. When you get to fancier stuff like joins, group-bys, etc., the mapping is not trivial at all.

#2

What about
“*real-time*”
event processing?
?

Wednesday, November 14, 12

Even Google now does not want to wait to compute a new index of the web. As soon as its web crawlers detect changes, Google would like to update its index as soon as possible.

The MapReduce Java API

39

Wednesday, November 14, 12

A Java API, but I'll show you Scala code; see my GitHub [scala-hadoop](#) project.

The MapReduce Java API (Problem #1)

39

Wednesday, November 14, 12

A Java API, but I'll show you Scala code; see my GitHub [scala-hadoop](#) project.

```

import org.apache.hadoop.io._
import org.apache.hadoop.mapred._
import java.util.StringTokenizer

object SimpleWordCount {

    val one = new IntWritable(1)
    val word = new Text      // Value will be set in a non-thread-safe way!

    class WCMapper extends MapReduceBase with Mapper[LongWritable, Text, Text, IntWritable] {

        def map(key: LongWritable, valueDocContents: Text,
               output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {
            val tokens = valueDocContents.toString.split("\\s+")
            for (wordString <- tokens) {
                if (wordString.length > 0) {
                    word.set(wordString.toLowerCase)
                    output.collect(word, one)
                }
            }
        }
    }

    class Reduce extends MapReduceBase with Reducer[Text, IntWritable, Text, IntWritable] {

        def reduce(keyWord: Text, valuesCounts: java.util.Iterator[IntWritable],
                   output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {
            var totalCount = 0
            while (valuesCounts.hasNext) {
                totalCount += valuesCounts.next.get
            }
            output.collect(keyWord, new IntWritable(totalCount))
        }
    }
}

```

This is intentionally too small to read. The algorithm is simple, but the framework is in your face. In the next several slides, notice which colors dominate. In this slide, it's green for types (classes), with relatively few yellow functions that implement actual operations.

Still, I've omitted many boilerplate details for configuring and running the job. This is just the "core" MapReduce code. In fact, Word Count is not too bad, but when you get to more complex algorithms, even conceptually simple ideas like relational-style joins and group-bys, the corresponding MapReduce code in this API gets complex and tedious very fast!



41

Wednesday, November 14, 12

Cascading is a Java library that provides higher-level abstractions for building data processing pipelines with concepts familiar from SQL such as a joins, group-bys, etc. It works on top of Hadoop's MapReduce and hides the boilerplate from you.

See <http://cascading.org>.

Use Cascading (Java)

(Solution #1a)

41

Wednesday, November 14, 12

Cascading is a Java library that provides higher-level abstractions for building data processing pipelines with concepts familiar from SQL such as a joins, group-bys, etc. It works on top of Hadoop's MapReduce and hides the boilerplate from you.

See <http://cascading.org>.

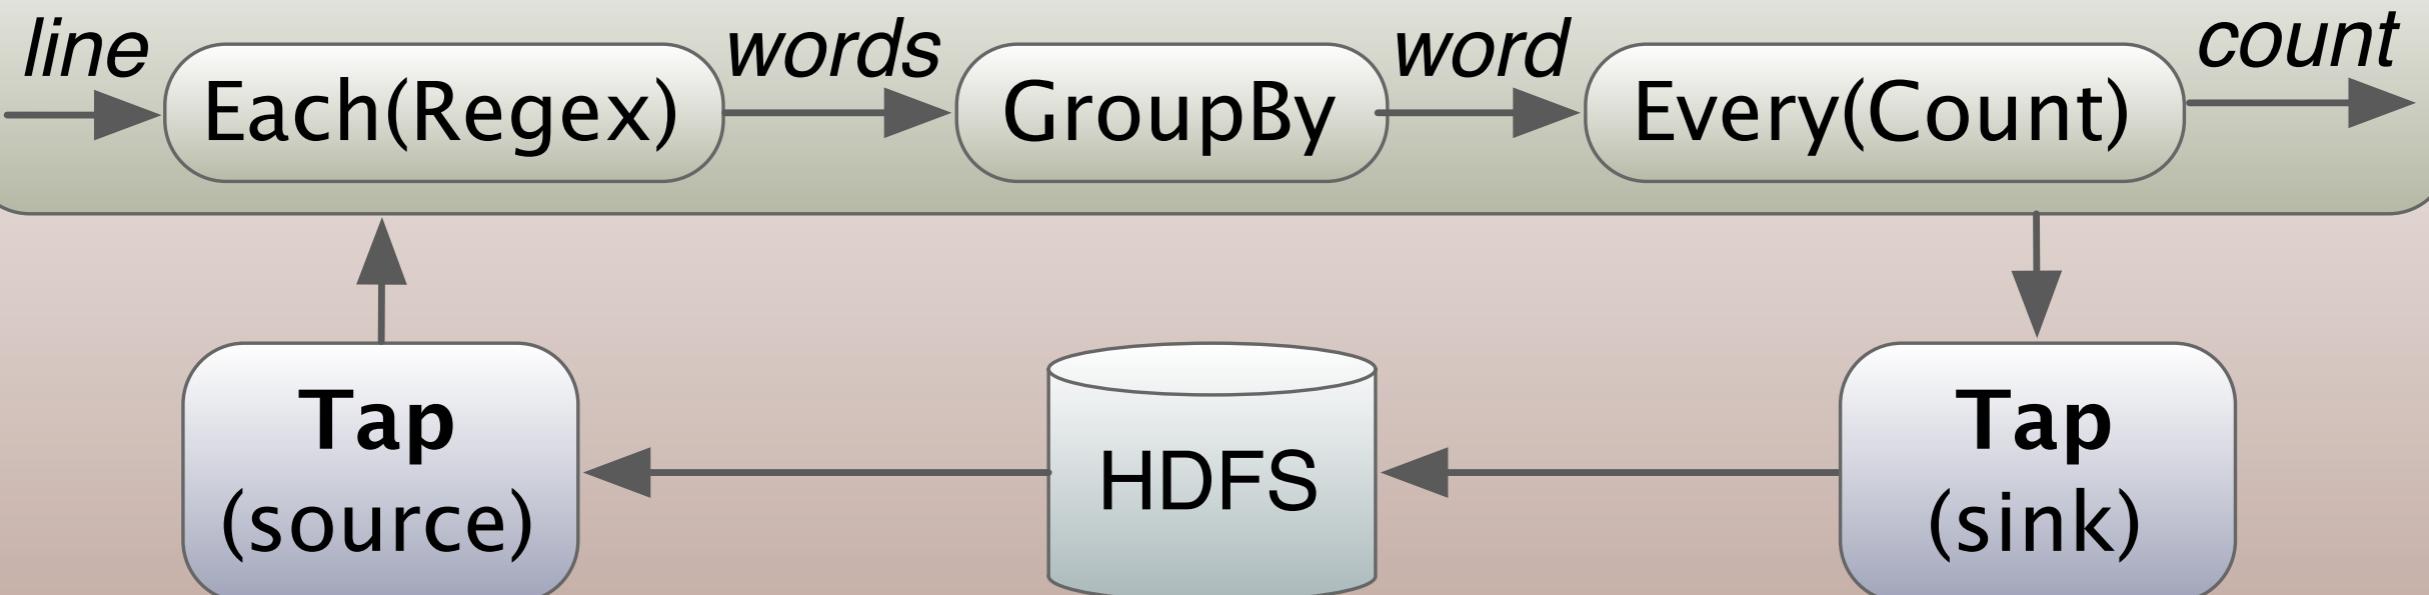
Cascading Concepts

Data flows consist of source and sink Taps connected by Pipes.

Word Count

Flow

Pipe ("word count assembly")



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Schematically, here is what Word Count looks like in Cascading. See <http://docs.cascading.org/cascading/1.2/userguide/html/ch02.html> for details.

```

import org.cascading.*;
...
public class WordCount {
    public static void main(String[] args) {
        String inputPath = args[0];
        String outputPath = args[1];
        Properties properties = new Properties();
        FlowConnector.setApplicationJarClass( properties, Main.class );

        Scheme sourceScheme = new TextLine( new Fields( "line" ) );
        Scheme sinkScheme = new TextLine( new Fields( "word", "count" ) );
        Tap source = new Hfs( sourceScheme, inputPath );
        Tap sink = new Hfs( sinkScheme, outputPath, SinkMode.REPLACE );

        Pipe assembly = new Pipe( "wordcount" );

        String regex = "(?<!\\pL)(?=\\pL)[^ ]*(?=<\\pL)(?!\\pL)";
        Function function = new RegexGenerator( new Fields( "word" ), regex );
        assembly = new Each( assembly, new Fields( "line" ), function );
        assembly = new GroupBy( assembly, new Fields( "word" ) );
        Aggregator count = new Count( new Fields( "count" ) );
        assembly = new Every( assembly, count );

        FlowConnector flowConnector = new FlowConnector( properties );
        Flow flow = flowConnector.connect( "word-count", source, sink, assembly );
        flow.complete();
    }
}

```

Here is the Cascading Java code. It's cleaner than the MapReduce API, because the code is more focused on the algorithm with less boilerplate, although it's not that much shorter. See <http://docs.cascading.org/cascading/1.2/userguide/html/ch02.html> for details; we won't discuss them here, but just go over highlights.

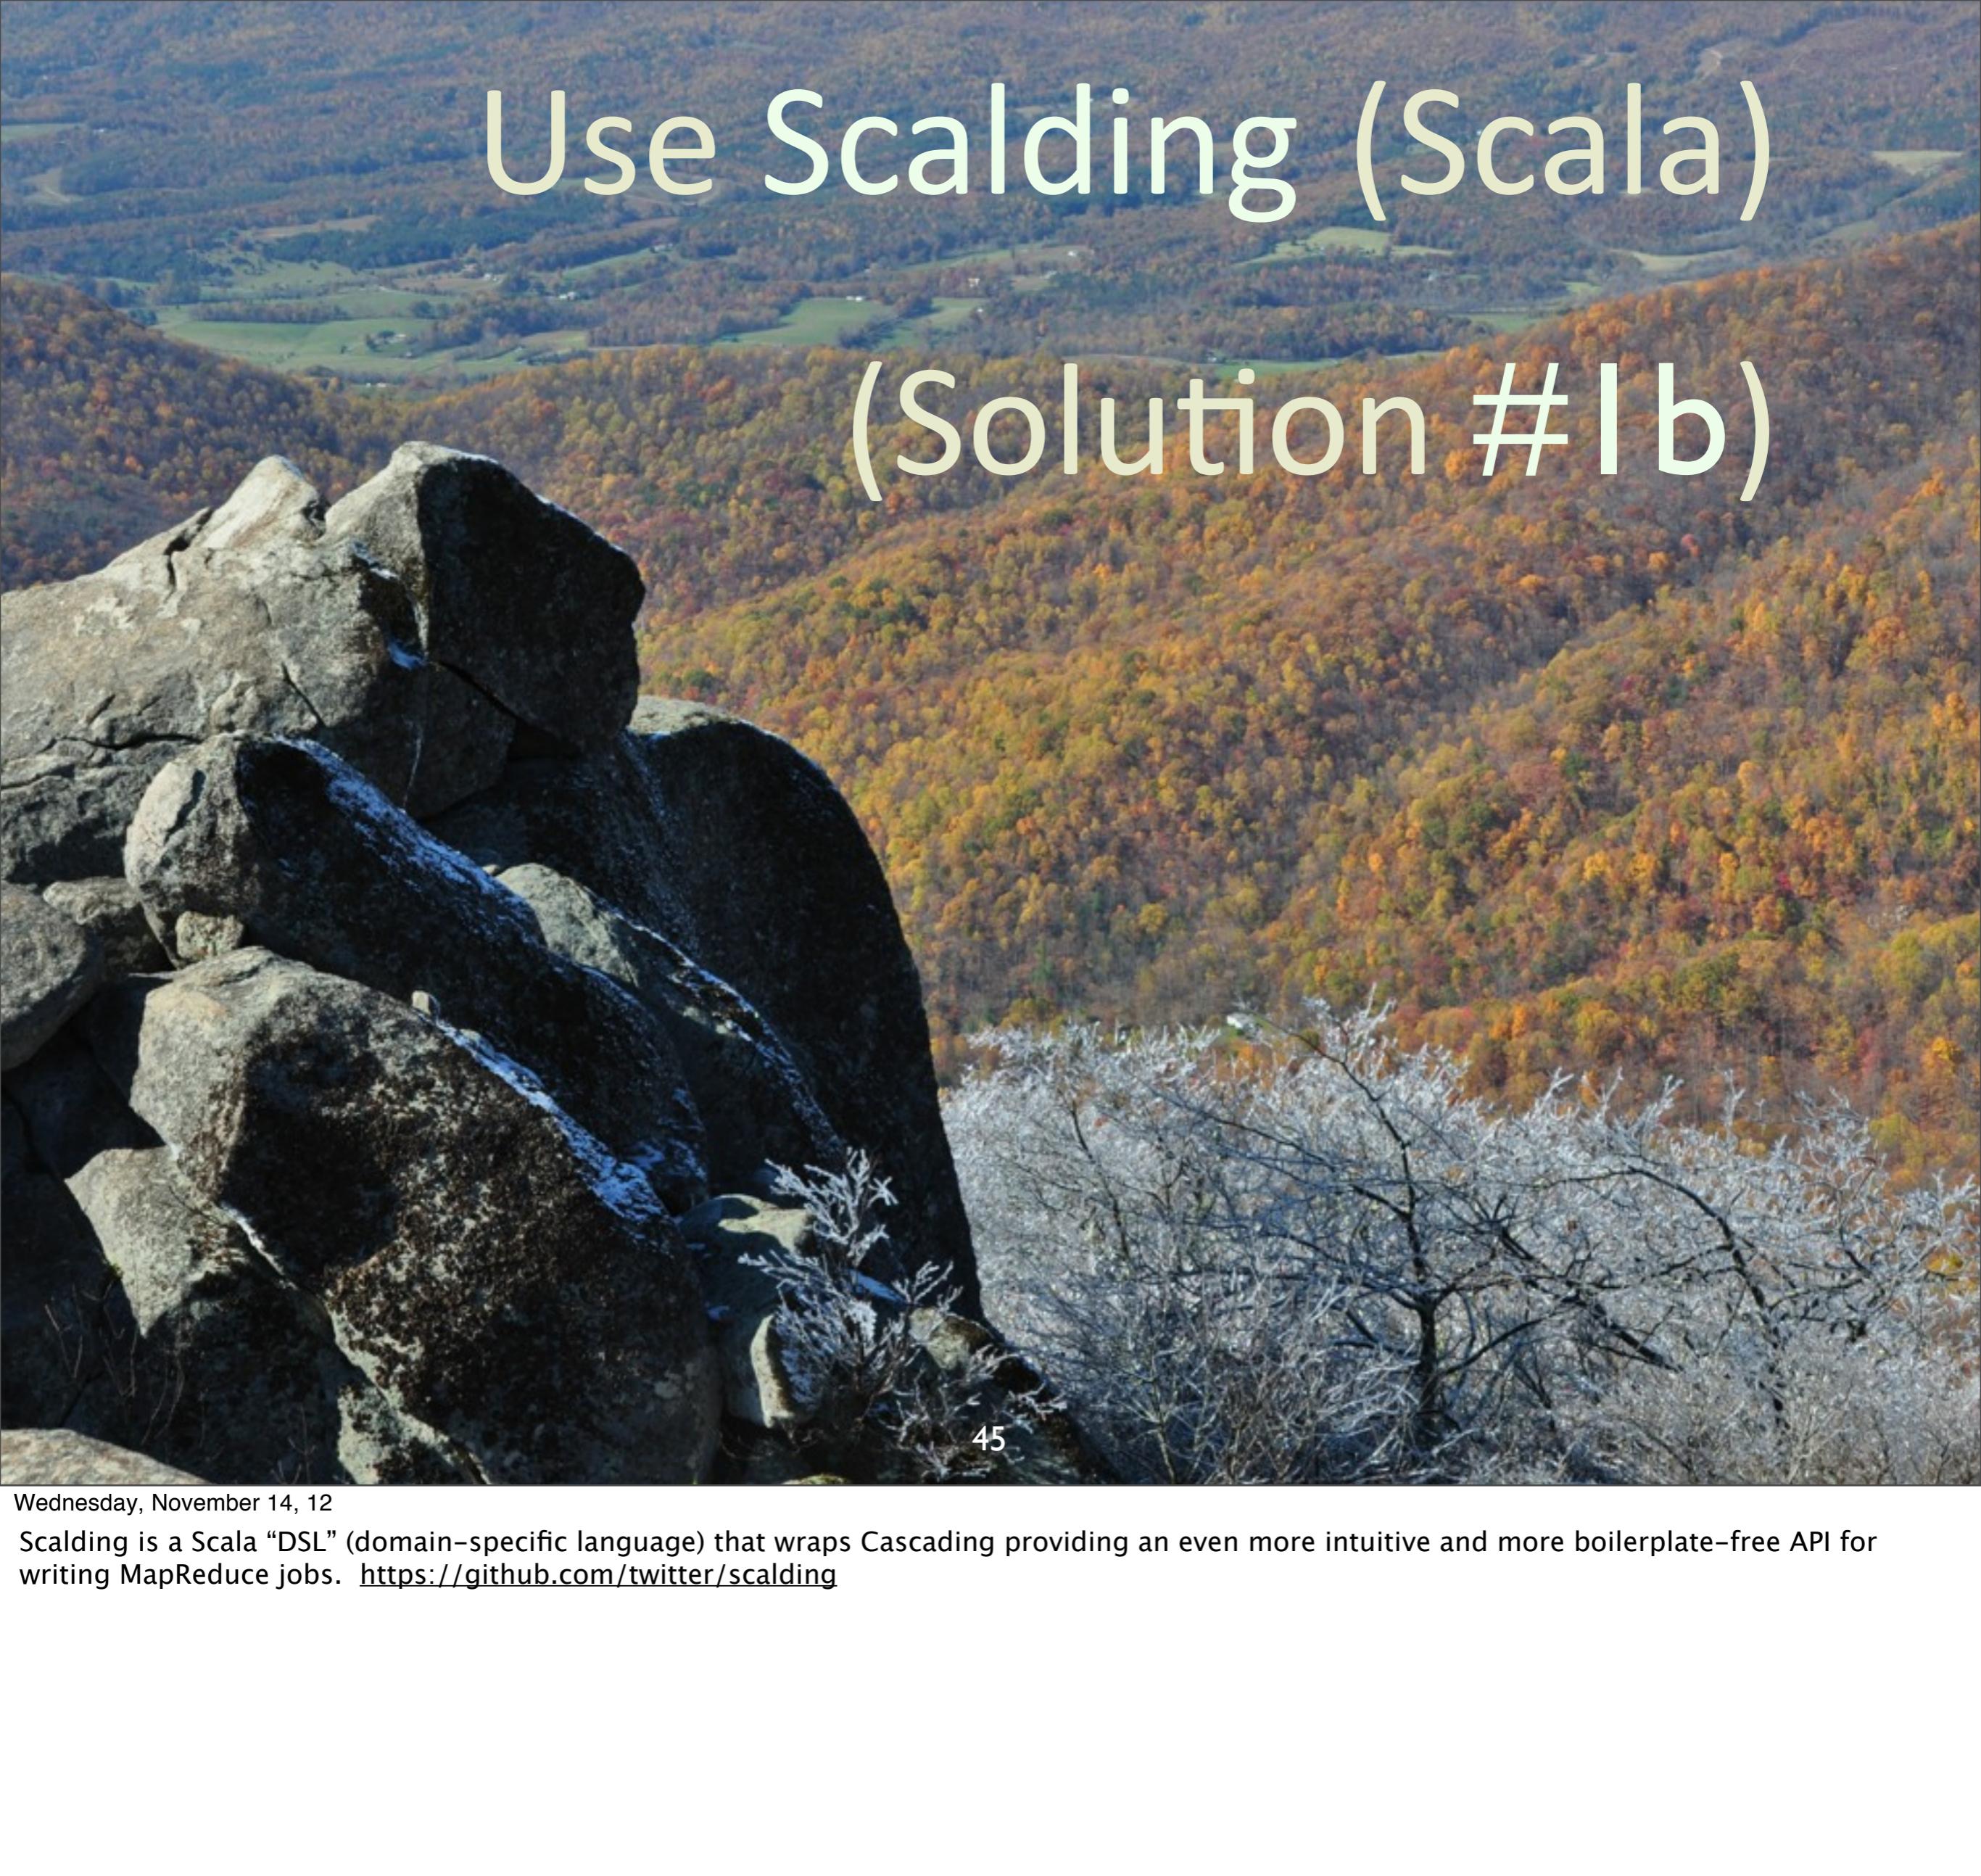
Note that there is still a lot of green for types, but at least the API emphasizes composing behaviors together.



45

Wednesday, November 14, 12

Scalding is a Scala “DSL” (domain-specific language) that wraps Cascading providing an even more intuitive and more boilerplate-free API for writing MapReduce jobs. <https://github.com/twitter/scalding>



Use Scalding (Scala) (Solution #1b)

45

Wednesday, November 14, 12

Scalding is a Scala “DSL” (domain-specific language) that wraps Cascading providing an even more intuitive and more boilerplate-free API for writing MapReduce jobs. <https://github.com/twitter/scalding>

```
import com.twitter.scalding._

class WordCountJob(args: Args) extends Job(args) {
  TextLine( args("input") )
    .read
    .flatMap('line -> 'word) {
      line: String => line.trim.toLowerCase.split("\\\\w+")
    }
    .groupBy('word) { group => group.size('count) }
  }
  .write(Tsv(args("output")))
}
```

This Scala code is almost pure domain logic with very little boilerplate. There are a few minor differences in the implementation. You don't explicitly specify the "Hfs" (Hadoop Distributed File System) taps. That's handled by Scalding implicitly when you run in "non-local" mode. Also, I'm using a simpler tokenization approach here, where I split on anything that isn't a "word character" [0-9a-zA-Z_].

There is little green, in part because Scala infers type in many cases. There is a lot more yellow for the methods that do real work!

What if MapReduce, and hence Cascading and Scalding, went obsolete tomorrow? This code is so short, I wouldn't care about throwing it away! I invested little time writing it, testing it, etc.

```
import com.twitter.scalding._

class WordCountJob(args: Args) extends Job(args) {
  TextLine( args("input") )
    .read
    .flatMap('line -> 'word) {
      line: String => line.trim.toLowerCase.split("\\\\w+")
    }
    .groupBy('word) { group => group.size('count) }
  }
  .write(Tsv(args("output")))
}
```

That's It!!

46

Wednesday, November 14, 12

This Scala code is almost pure domain logic with very little boilerplate. There are a few minor differences in the implementation. You don't explicitly specify the "Hfs" (Hadoop Distributed File System) taps. That's handled by Scalding implicitly when you run in "non-local" model. Also, I'm using a simpler tokenization approach here, where I split on anything that isn't a "word character" [0-9a-zA-Z_].

There is little green, in part because Scala infers type in many cases. There is a lot more yellow for the methods that do real work!

What if MapReduce, and hence Cascading and Scalding, went obsolete tomorrow? This code is so short, I wouldn't care about throwing it away! I invested little time writing it, testing it, etc.

You also see this
functional improvement
with other APIs:

- Crunch (Java)
- Scrunch (Scala)

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

See <https://github.com/cloudera/crunch>



48

Wednesday, November 14, 12

Scalding is a Scala “DSL” (domain-specific language) that wraps Cascading providing an even more intuitive and more boilerplate-free API for writing MapReduce jobs. <https://github.com/twitter/scalding>



Use Spark (Scala)

(Solution #2)

48

Wednesday, November 14, 12

Scalding is a Scala “DSL” (domain-specific language) that wraps Cascading providing an even more intuitive and more boilerplate-free API for writing MapReduce jobs. <https://github.com/twitter/scalding>

Spark is a Hadoop MapReduce alternative:

- Distributed computing with in-memory caching.
- Up to 30x faster than MapReduce.

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

See <http://www.spark-project.org/>

Why isn't it more widely used; lack of commercial support...

Spark is a Hadoop MapReduce alternative:

- Designed for machine learning applications.

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

See <http://www.spark-project.org/>

```
object WordCountSpark {  
    def main(args: Array[String]) {  
        val file = spark.textFile(args(0))  
        val counts = file.flatMap(line => line.split("\\\\w+"))  
                      .map(word => (word, 1))  
                      .reduceByKey(_ + _)  
        counts.saveAsTextFile(args(1))  
    }  
}
```

```
object WordCountSpark {  
    def main(args: Array[String]) {  
        val file = spark.textFile(args(0))  
        val counts = file.flatMap(line => line.split("\\\\w+"))  
                      .map(word => (word, 1))  
                      .reduceByKey(_ + _)  
        counts.saveAsTextFile(args(1))  
    }  
}
```

Also that's it!
Note it's similar to the MapReduce API,
but far more concise.



52

Wednesday, November 14, 12

Using SQL when you can! Here are 3 options.

Use Hive, Shark, or Impala

(Solution #3)

Use SQL when you can:

- Hive: SQL on top of MapReduce.
- Shark: Hive ported to Spark.
- Impala: Hive SQL with new, faster back end.

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

See <http://hive.apache.org/> or my book for Hive, <http://shark.cs.berkeley.edu/> for shark, and <http://www.cloudera.com/content/cloudera/en/products/cloudera-enterprise-core/cloudera-enterprise-RTQ.html> for Impala.

SQL!

```
CREATE TABLE docs (line STRING);
LOAD DATA INPATH '/path/to/docs' INTO TABLE docs;
```

```
CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
(SELECT explode(split(line, '\W+')) AS word FROM docs) w
GROUP BY word
ORDER BY word;
```

SQL!

```
CREATE TABLE docs (line STRING);
LOAD DATA INPATH '/path/to/docs' INTO TABLE docs;
```

```
CREATE TABLE word_counts AS
SELECT word, count(1) AS count FROM
(SELECT explode(split(line, '\W+')) AS word FROM docs) w
GROUP BY word
ORDER BY word;
```

Word Count, again...
... in HiveQL

54

Wednesday, November 14, 12

This is how you could implement word count in Hive. We're using some Hive built-in functions for tokenizing words in each "line", the one "column" in the docs table.

Impala

- HiveQL front end.
- C++ and Java back end.
- Provides up to 100x performance improvement!
- Developed by Cloudera.

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

See <http://www.cloudera.com/content/cloudera/en/products/cloudera-enterprise-core/cloudera-enterprise-RTQ.html>.



56

Wednesday, November 14, 12

A good summary presentation: <http://www.slideshare.net/shatteredNirvana/pregel-a-system-for-largescale-graph-processing>

Use Graph Processing (Solution #4)

56

Wednesday, November 14, 12

A good summary presentation: <http://www.slideshare.net/shatteredNirvana/pregel-a-system-for-largescale-graph-processing>

Google's Page Rank

- Google invented MapReduce,
- ... but MapReduce is suboptimal for Page Rank and other graph algorithms.

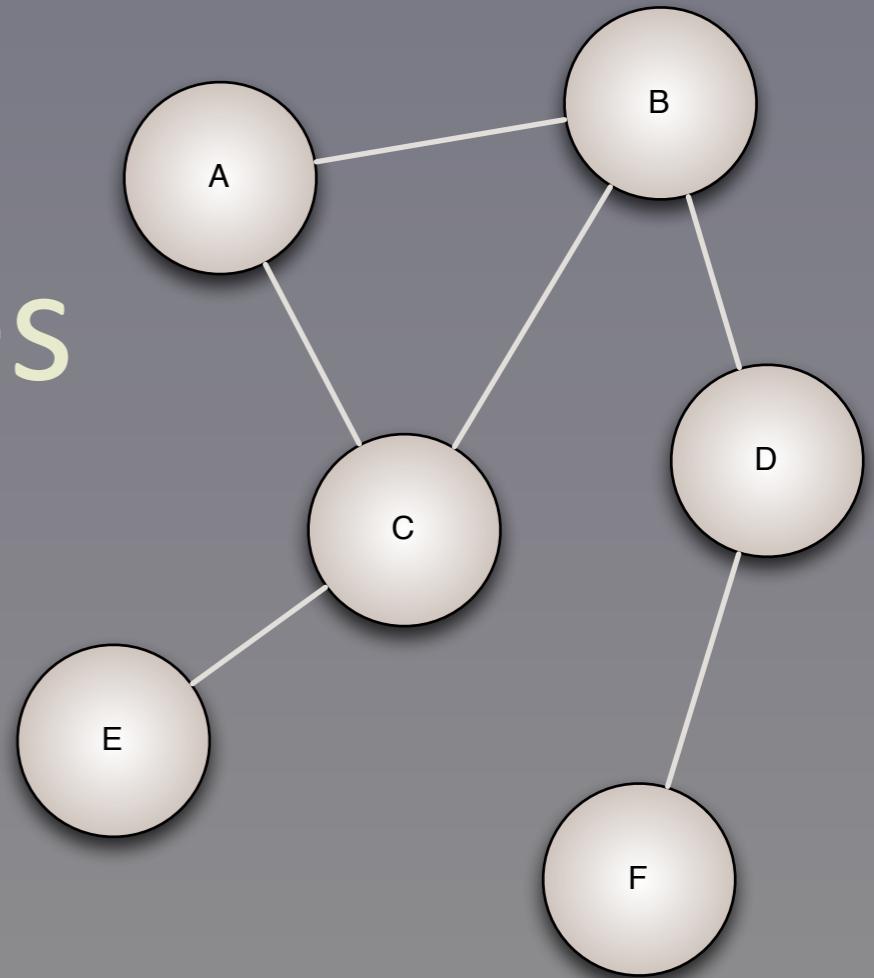
Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

PageRank is the famous algorithm invented by Sergey Brin and Larry Page to index the web. It's the foundation of Google's search engine.

Why not MapReduce?

- 1 MR job for each iteration that updates all n nodes/edges.
- Graph saved to disk after each iteration.
- ...



Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

The presentation <http://www.slideshare.net/shatteredNirvana/pregel-a-system-for-largescale-graph-processing> itemizes all the major issues with using MR to implement graph algorithms.

Google's Pregel

- Page Rank now runs on a new graph framework: Pregel.
- Bulk, Synchronous Parallel (BSP).
 - Graphs are first-class citizens.
 - Efficiently processes updates...

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

Pregel is the name of the river that runs through the city of Königsberg, Prussia (now called Kaliningrad, Ukraine). 7 bridges crossed the river in the city (including to 5 to 2 islands between river branches). Leonhard Euler invented graph theory when we analyzed the question of whether or not you can cross all 7 bridges without retracing your steps (you can't).

Open-source Alternatives

- Apache Giraph.
- Apache Hama.
- Aurelius Titan.

Copyright © 2011-2012, Think Big Analytics, All Rights Reserved

Wednesday, November 14, 12

<http://incubator.apache.org/giraph/>

<http://hama.apache.org/>

<http://thinkaurelius.github.com/titan/>

None is very mature nor has commercial support.



61

Wednesday, November 14, 12

So, I think we have opportunities for a better way...



A Manifesto...

61

Wednesday, November 14, 12

So, I think we have opportunities for a better way...



Wednesday, November 14, 12

I worked with EJBs a decade ago. The framework was completely invasive into your business logic. There were too many configuration options in XML files. The framework “paradigm” was a poor fit for most problems (like soft real time systems and most algorithms beyond Word Count). Internally, EJB implementations were inefficient and hard to optimize, because they relied on poorly considered object boundaries that muddled more natural boundaries. (I’ve argued in other presentations and my “FP for Java Devs” book that OOP is a poor modularity tool...) The fact is, Hadoop reminds me of EJBs in almost every way. It works okay and people do get stuff done, but just as the Spring Framework brought an essential rethinking to Enterprise Java, I think there is an essential rethink that needs to happen in Big Data. The functional programming community, is well positioned to create it.



Hadoop is the Enterprise Java Beans of our time.

Wednesday, November 14, 12

I worked with EJBs a decade ago. The framework was completely invasive into your business logic. There were too many configuration options in XML files. The framework “paradigm” was a poor fit for most problems (like soft real time systems and most algorithms beyond Word Count). Internally, EJB implementations were inefficient and hard to optimize, because they relied on poorly considered object boundaries that muddled more natural boundaries. (I’ve argued in other presentations and my “FP for Java Devs” book that OOP is a poor modularity tool...) The fact is, Hadoop reminds me of EJBs in almost every way. It works okay and people do get stuff done, but just as the Spring Framework brought an essential rethinking to Enterprise Java, I think there is an essential rethink that needs to happen in Big Data. The functional programming community, is well positioned to create it.



Wednesday, November 14, 12

Java is really the wrong language for this and the continued use by Big Data vendors is slowing down progress. It offers the wrong abstractions, objects instead mathematically-inspired functional programming, and it's incredibly verbose when compared to modern alternatives.

Stop Using Java!



Wednesday, November 14, 12

Java is really the wrong language for this and the continued use by Big Data vendors is slowing down progress. It offers the wrong abstractions, objects instead mathematically-inspired functional programming, and it's incredibly verbose when compared to modern alternatives.



Wednesday, November 14, 12

Why is Functional Programming better for Big Data? The work we do with data is inherently mathematical transformations and FP is inspired by math. Hence, it's naturally a better fit, much more so than object-oriented programming. And, modern languages like Scala and Clojure are more concise and better at eliminating boilerplate.

SQL has succeeded all these years because it is also inspired by math, e.g., set theory.



Functional Languages improve Big Data productivity!

64

Wednesday, November 14, 12

Why is Functional Programming better for Big Data? The work we do with data is inherently mathematical transformations and FP is inspired by math. Hence, it's naturally a better fit, much more so than object-oriented programming. And, modern languages like Scala and Clojure are more concise and better at eliminating boilerplate.

SQL has succeeded all these years because it is also inspired by math, e.g., set theory.



Wednesday, November 14, 12

We already have the right model in the collection APIs that come with functional languages. They are far better engineered for intuitive data transformations. They provide the right abstractions and hide boilerplate. In fact, they make it relatively easy to optimization implementations for parallelization. The Scala collections offer parallelization with a tiny API call. Spark and Cascading transparently distribute collections across a cluster.



Functional Collections.

Wednesday, November 14, 12

We already have the right model in the collection APIs that come with functional languages. They are far better engineered for intuitive data transformations. They provide the right abstractions and hide boilerplate. In fact, they make it relatively easy to optimization implementations for parallelization. The Scala collections offer parallelization with a tiny API call. Spark and Cascading transparently distribute collections across a cluster.



Wednesday, November 14, 12

We can start using new, more efficient compute models, like Spark, Pregel, and Impala today. Of course, you have to consider maturity, viability, and support issues in large organizations. So if you want to wait until these alternatives are more mature, then at least use better APIs for Hadoop!

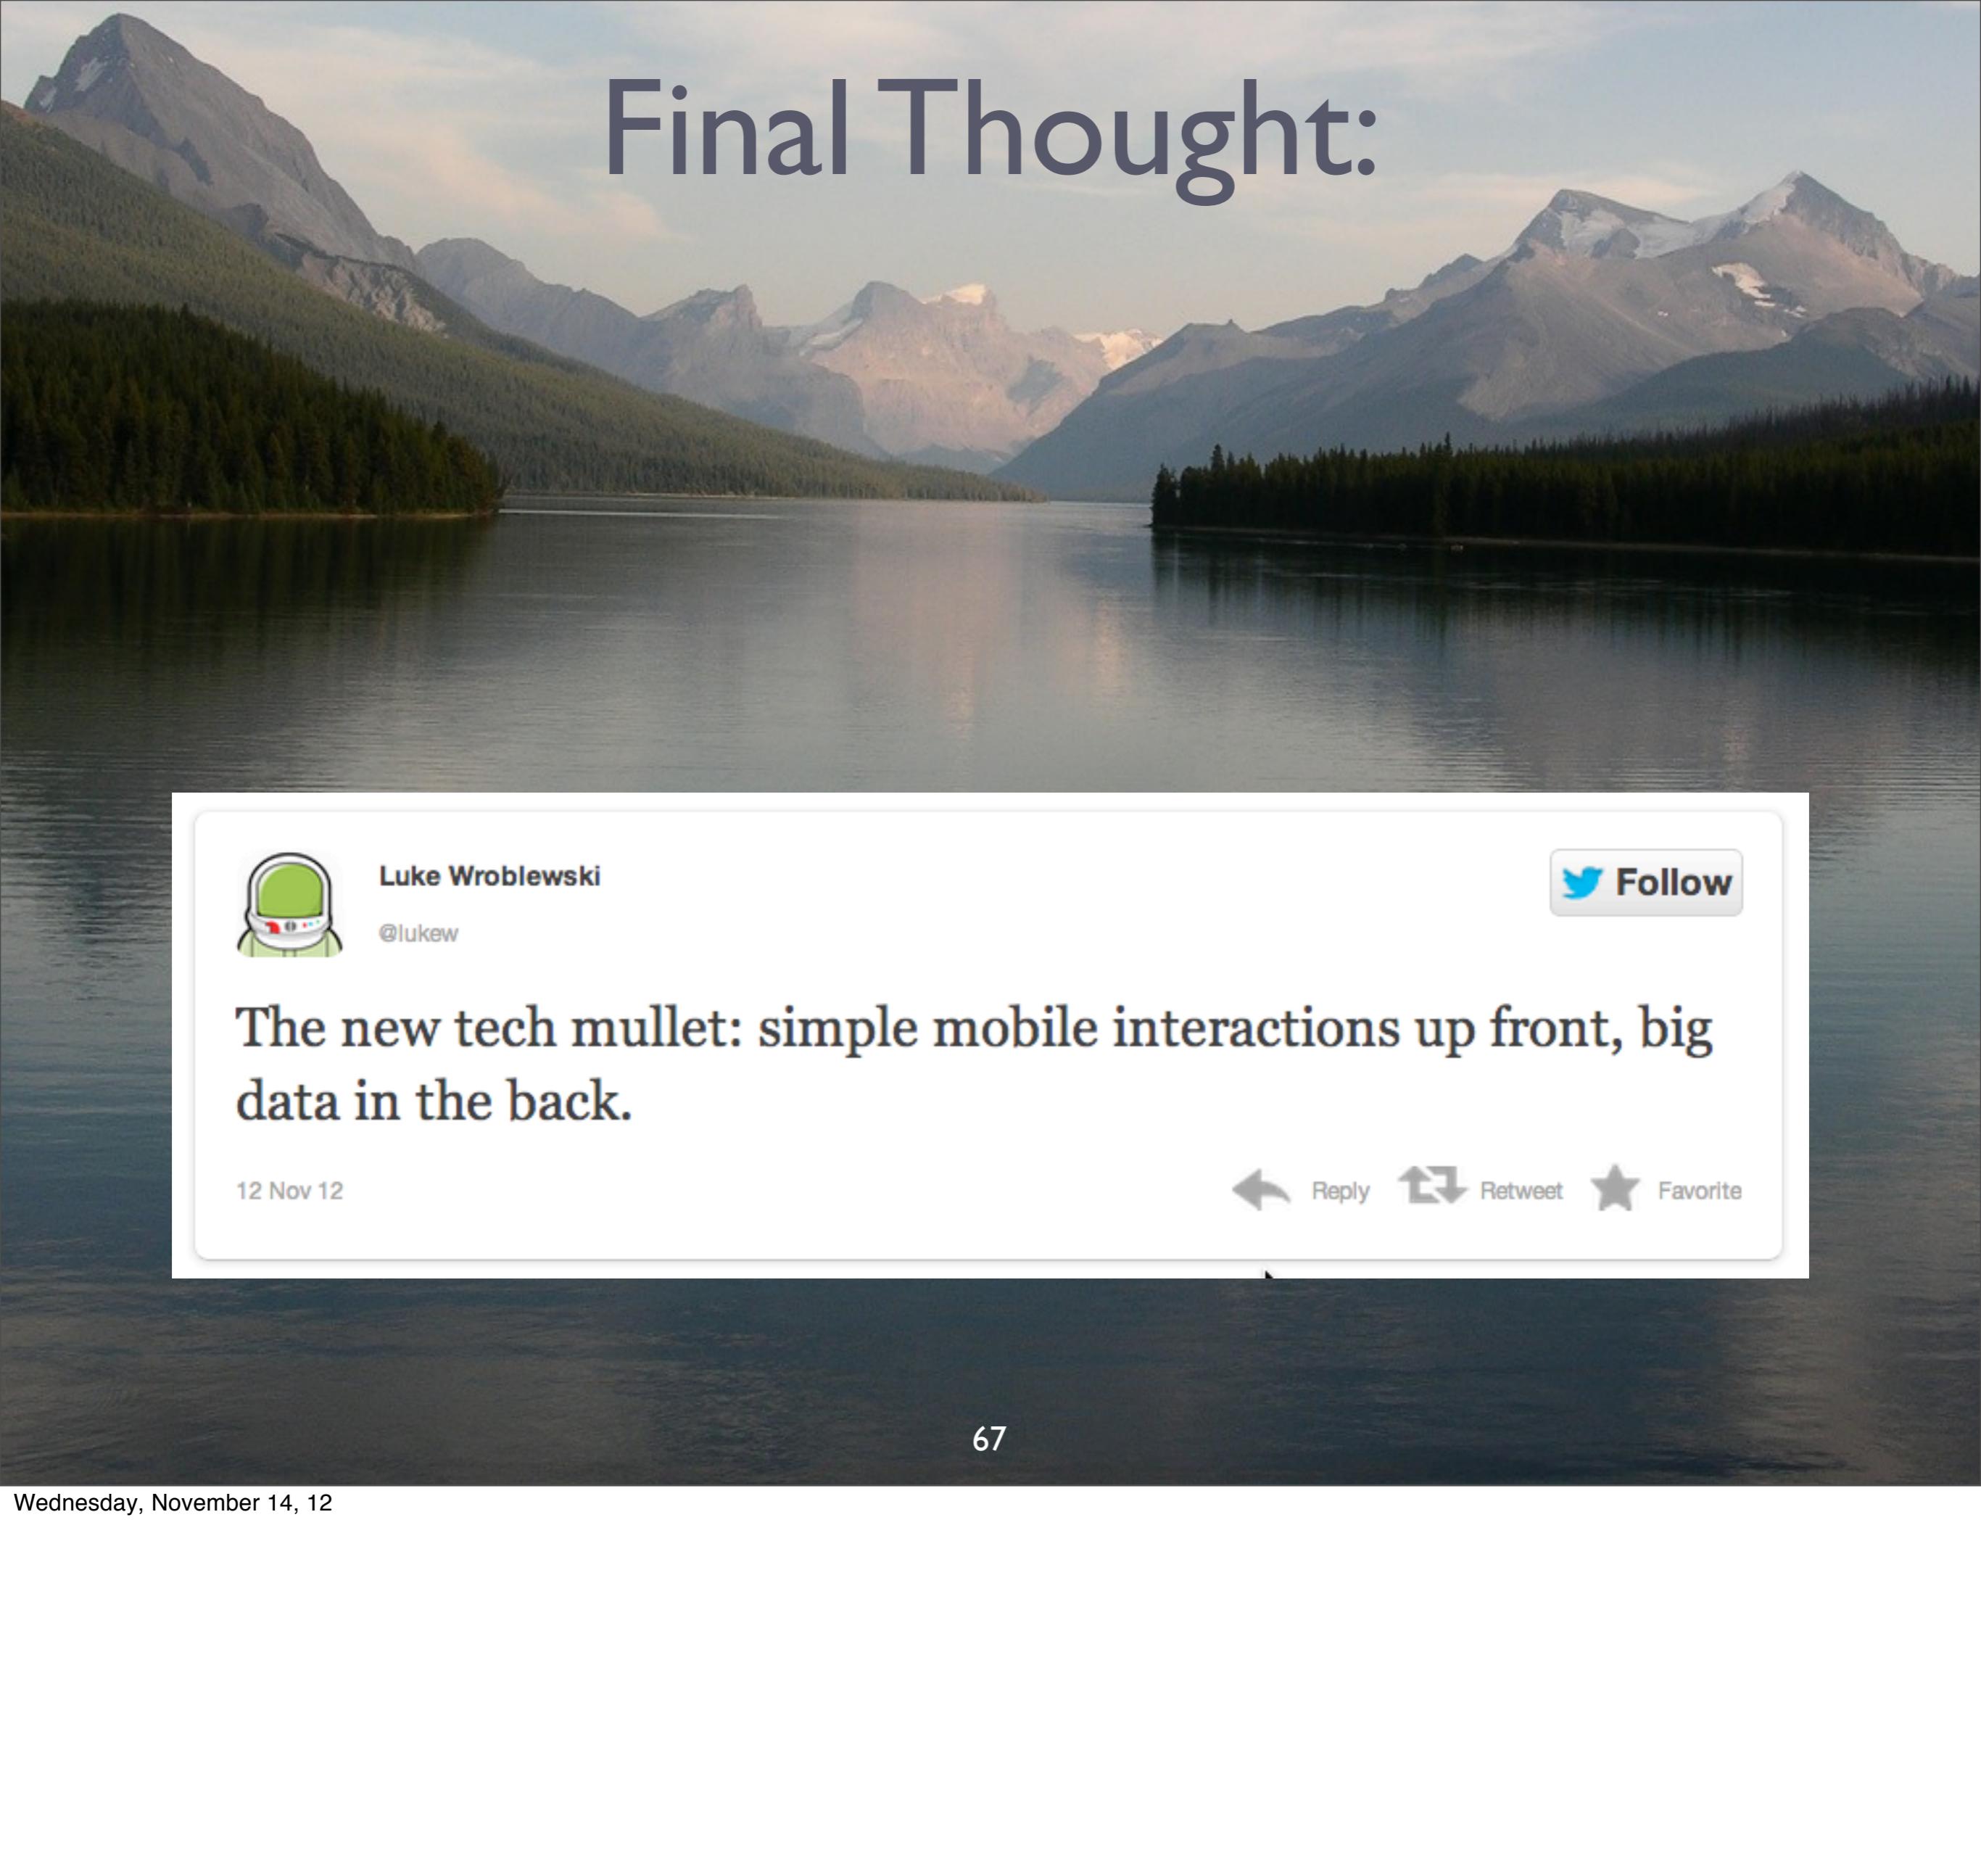
A scenic mountain landscape featuring a foreground of vibrant yellow wildflowers, likely sunflowers, stretching across the frame. Behind them, a dense forest of tall evergreen trees lines a valley. In the background, majestic mountains rise, their slopes covered with patches of snow and green vegetation. The sky above is a clear, bright blue, dotted with wispy white clouds.

New Compute Models.

Wednesday, November 14, 12

We can start using new, more efficient compute models, like Spark, Pregel, and Impala today. Of course, you have to consider maturity, viability, and support issues in large organizations. So if you want to wait until these alternatives are more mature, then at least use better APIs for Hadoop!

Final Thought:



Luke Wroblewski

@lukew

The new tech mullet: simple mobile interactions up front, big data in the back.

12 Nov 12

Follow

Reply Retweet Favorite

67

Wednesday, November 14, 12

Questions?

QCon Motorola Solutions SABA
November 13, 2012

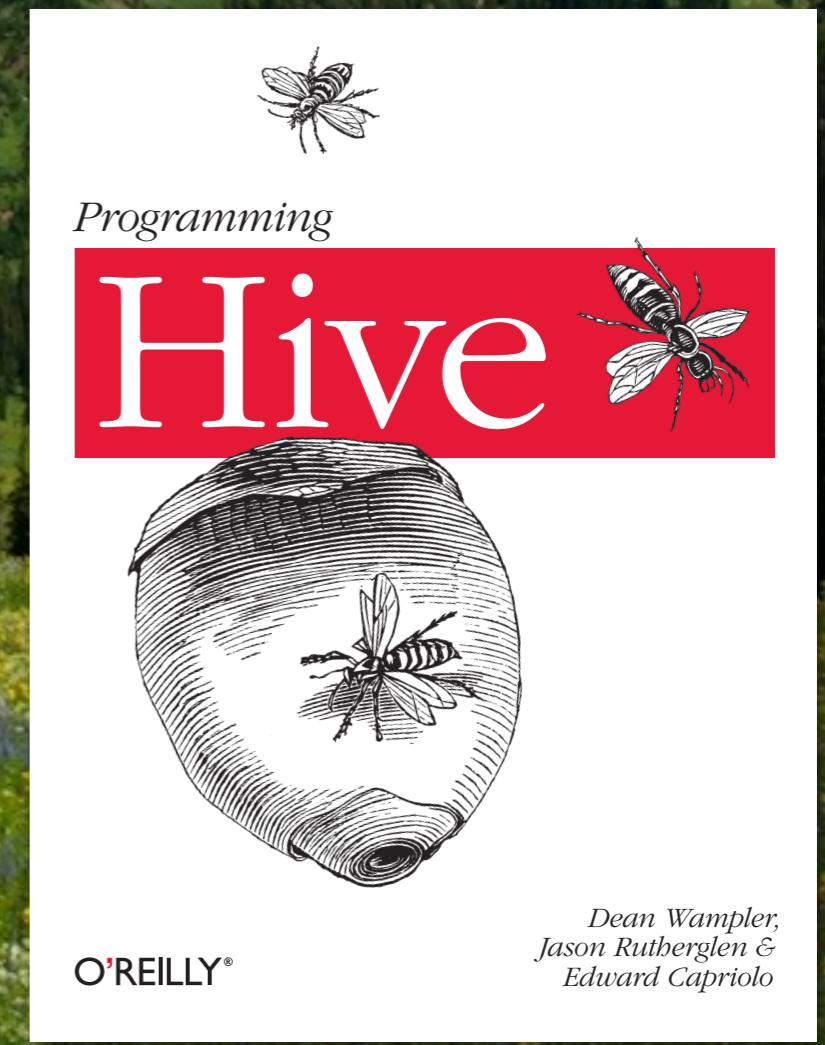
dean.wampler@thinkbiganalytics.com

68



Wednesday, November 14, 12

All pictures © Dean Wampler, 2011-2012.



Questions?

QCon Motorola Solutions SABA
November 13, 2012

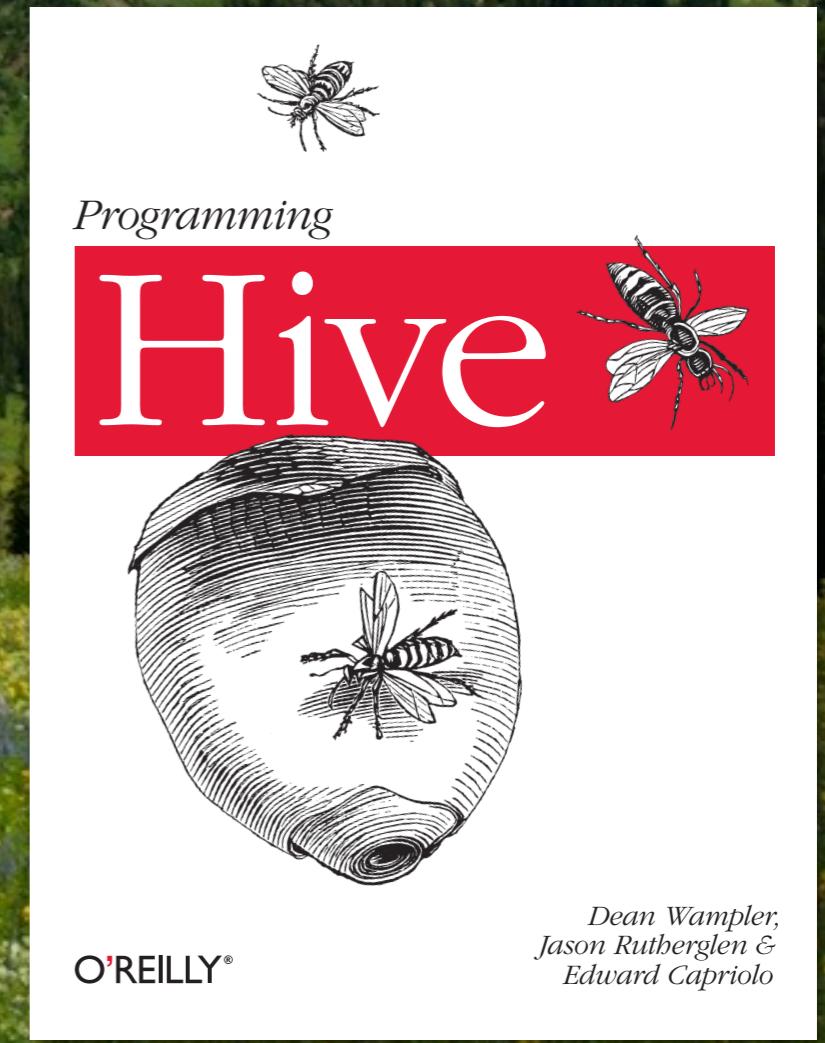
dean.wampler@thinkbiganalytics.com

68



Wednesday, November 14, 12

All pictures © Dean Wampler, 2011-2012.



SOME RIGHTS RESERVED