

QCon NYC
June 19, 2012
@deanwampler
thinkbiganalytics.com

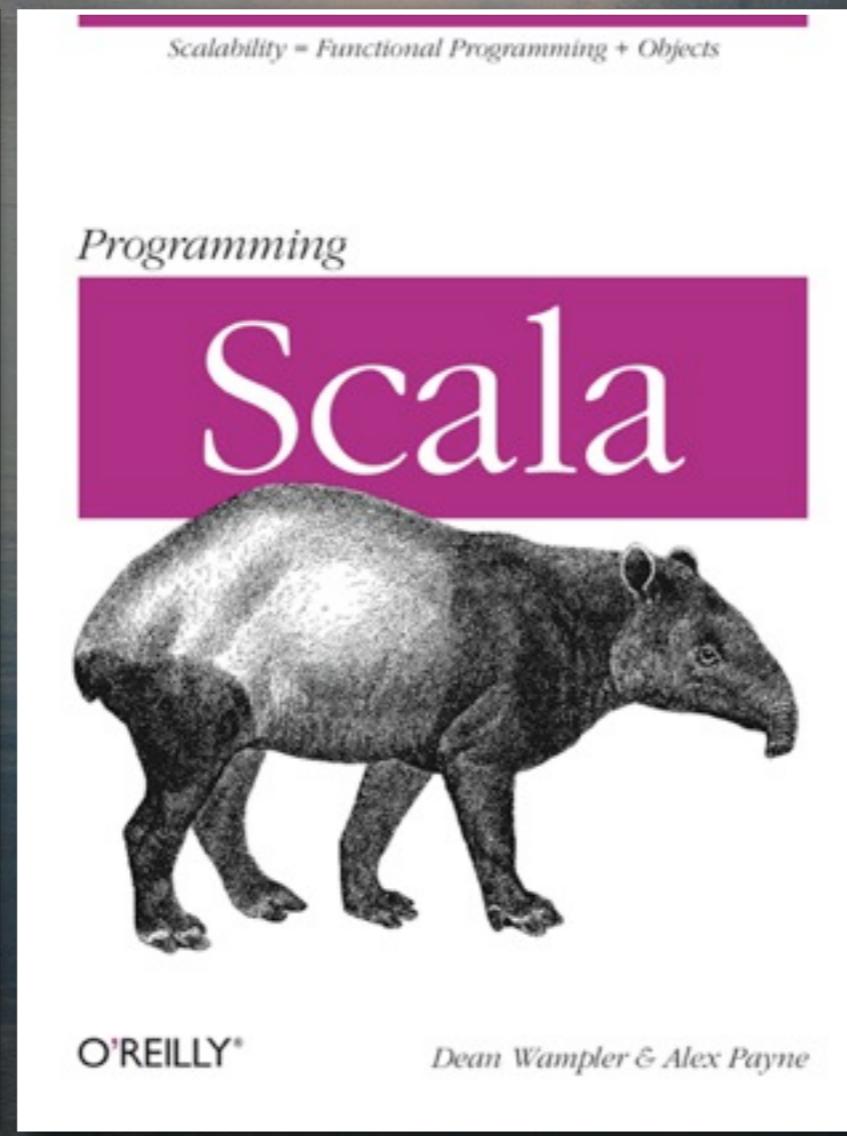
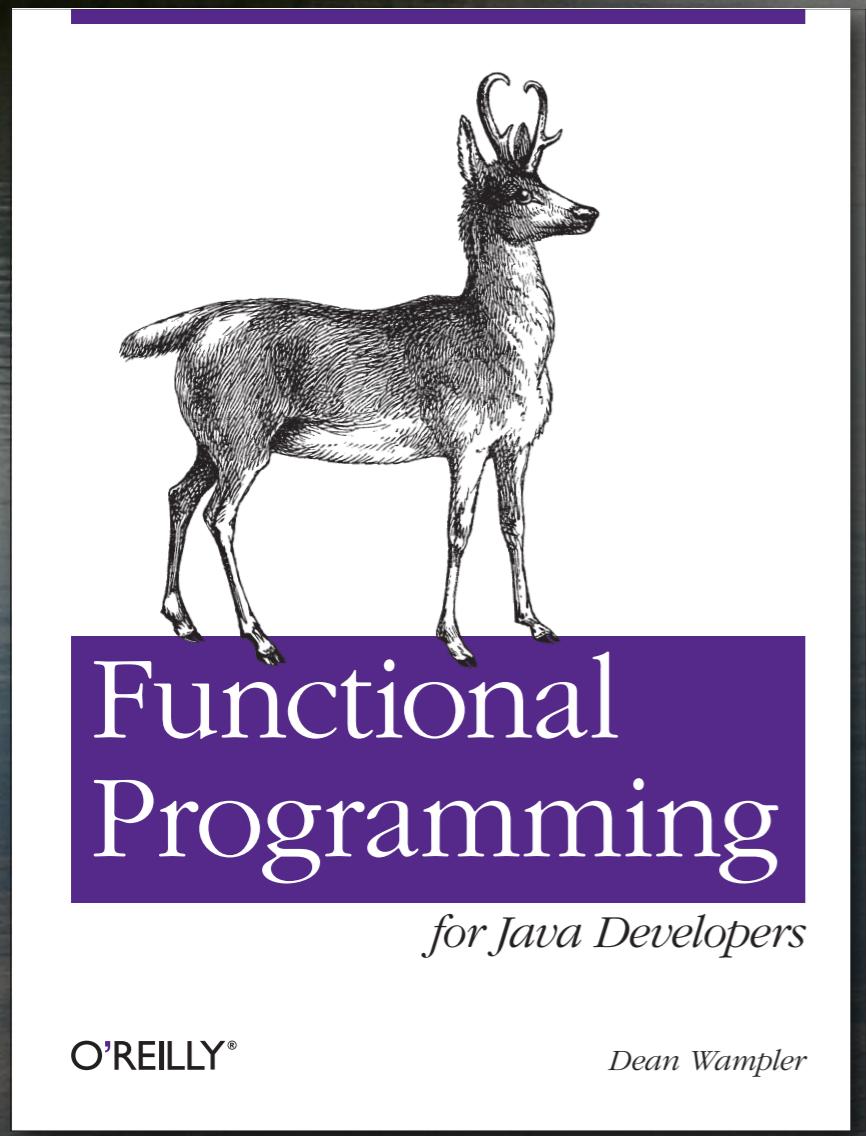
MapReduce and Its Discontents



Tuesday, June 19, 12

MapReduce is a powerful tool for scaling data analytics and storage. It is on the hype curve of popularity. However, is MR the end of the story? Does it meet all our needs?

<shameless/>



Tuesday, June 19, 12

"Programming Hive" is forthcoming, Sept. 2012.

Big Data

Data so big that traditional solutions are too slow, too small, or too expensive to use.



Hat tip: Bob Korbus

Tuesday, June 19, 12

It's a buzz word, but generally associated with the problem of data sets too big to manage with traditional SQL databases. A parallel development has been the NoSQL movement that is good at handling semistructured data, scaling, etc.



3 Trends

Tuesday, June 19, 12

Three trends influence my thinking...

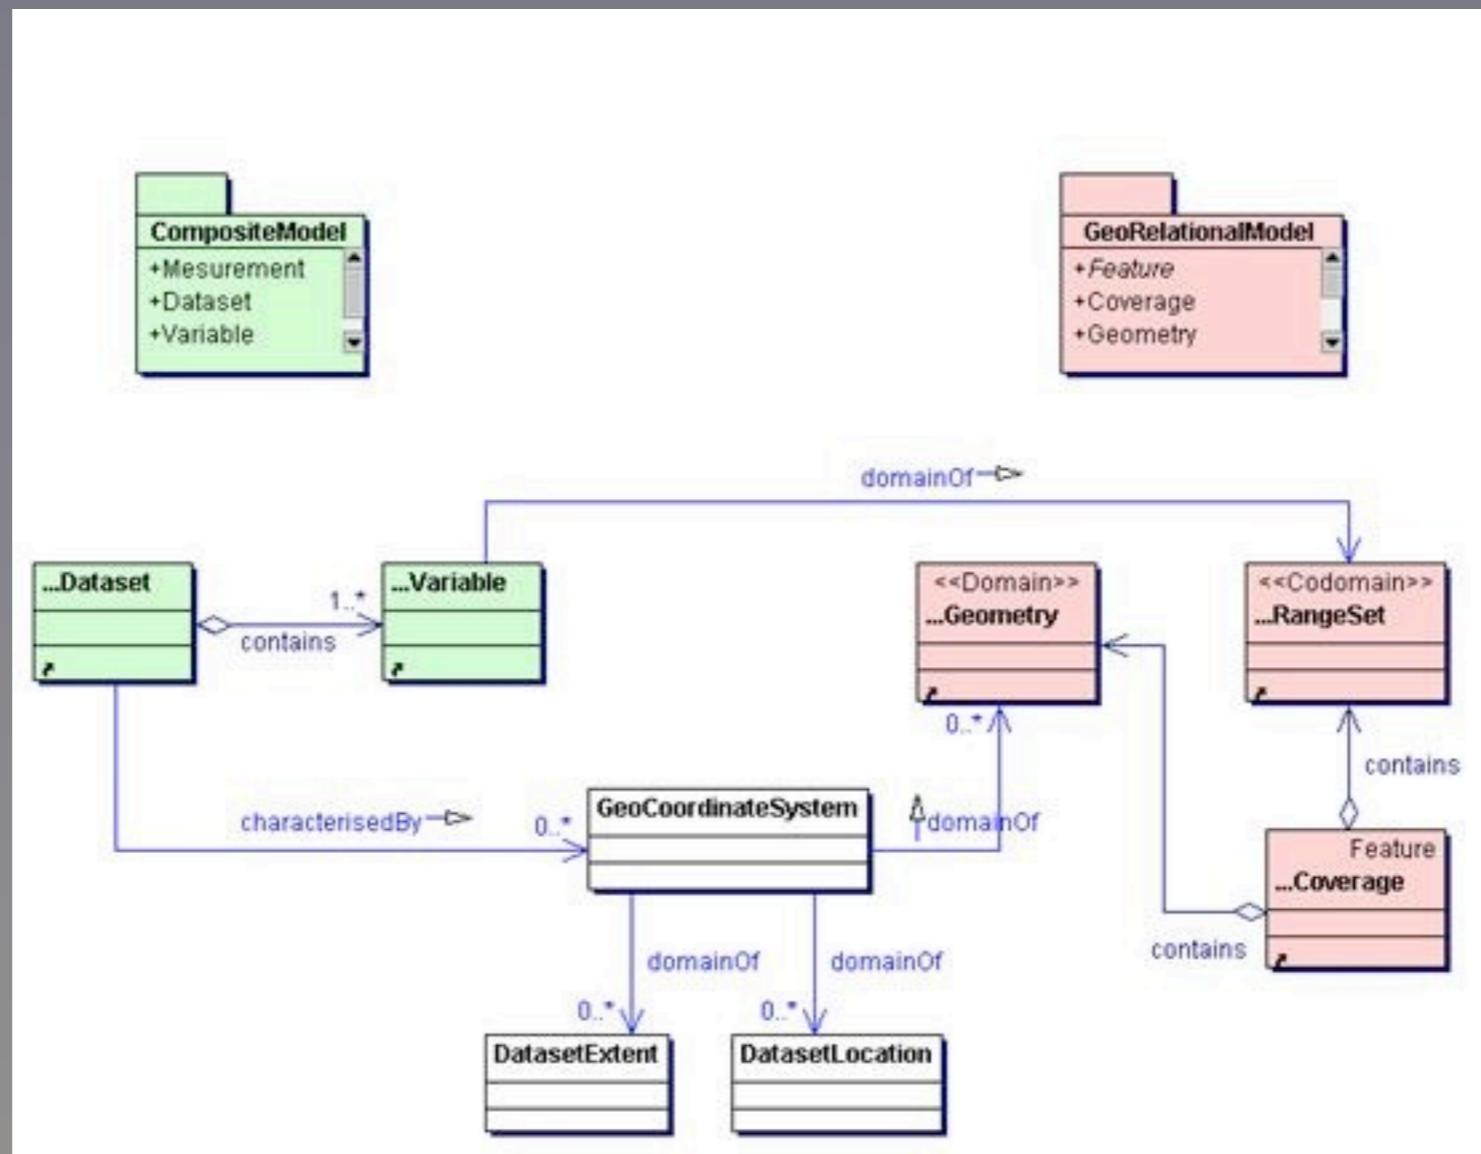
Data Size ↑



Tuesday, June 19, 12

Data volumes are obviously growing... rapidly.

Formal Schemas



Tuesday, June 19, 12

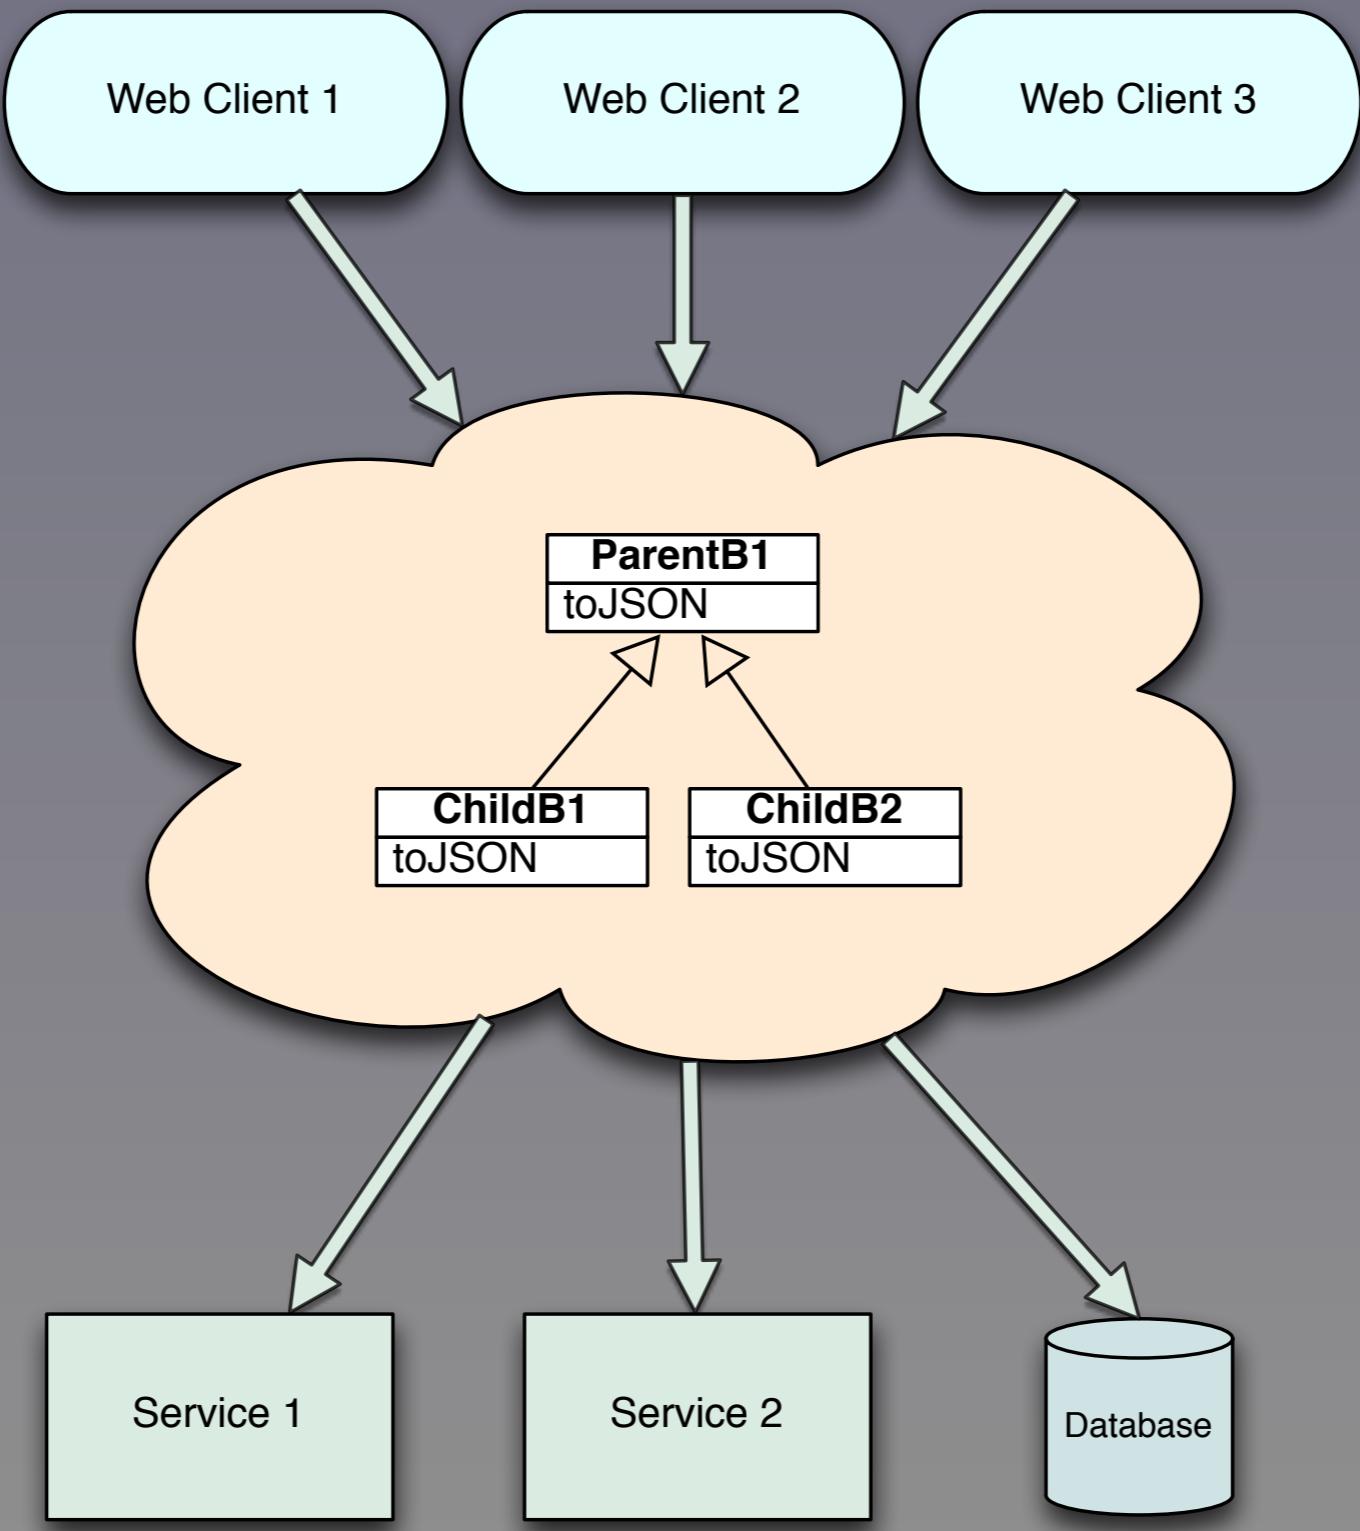
There is less emphasis on “formal” schemas and domain models, i.e., both relational models of data and OO models, because data schemas and sources change rapidly. So, using relatively-agnostic software (e.g., collections of things where the software is agnostic about the contents) tends to be faster to develop and run. Put another way, we find it more useful to build somewhat agnostic applications and drive their behavior through data.

Data-Driven Programs ↑



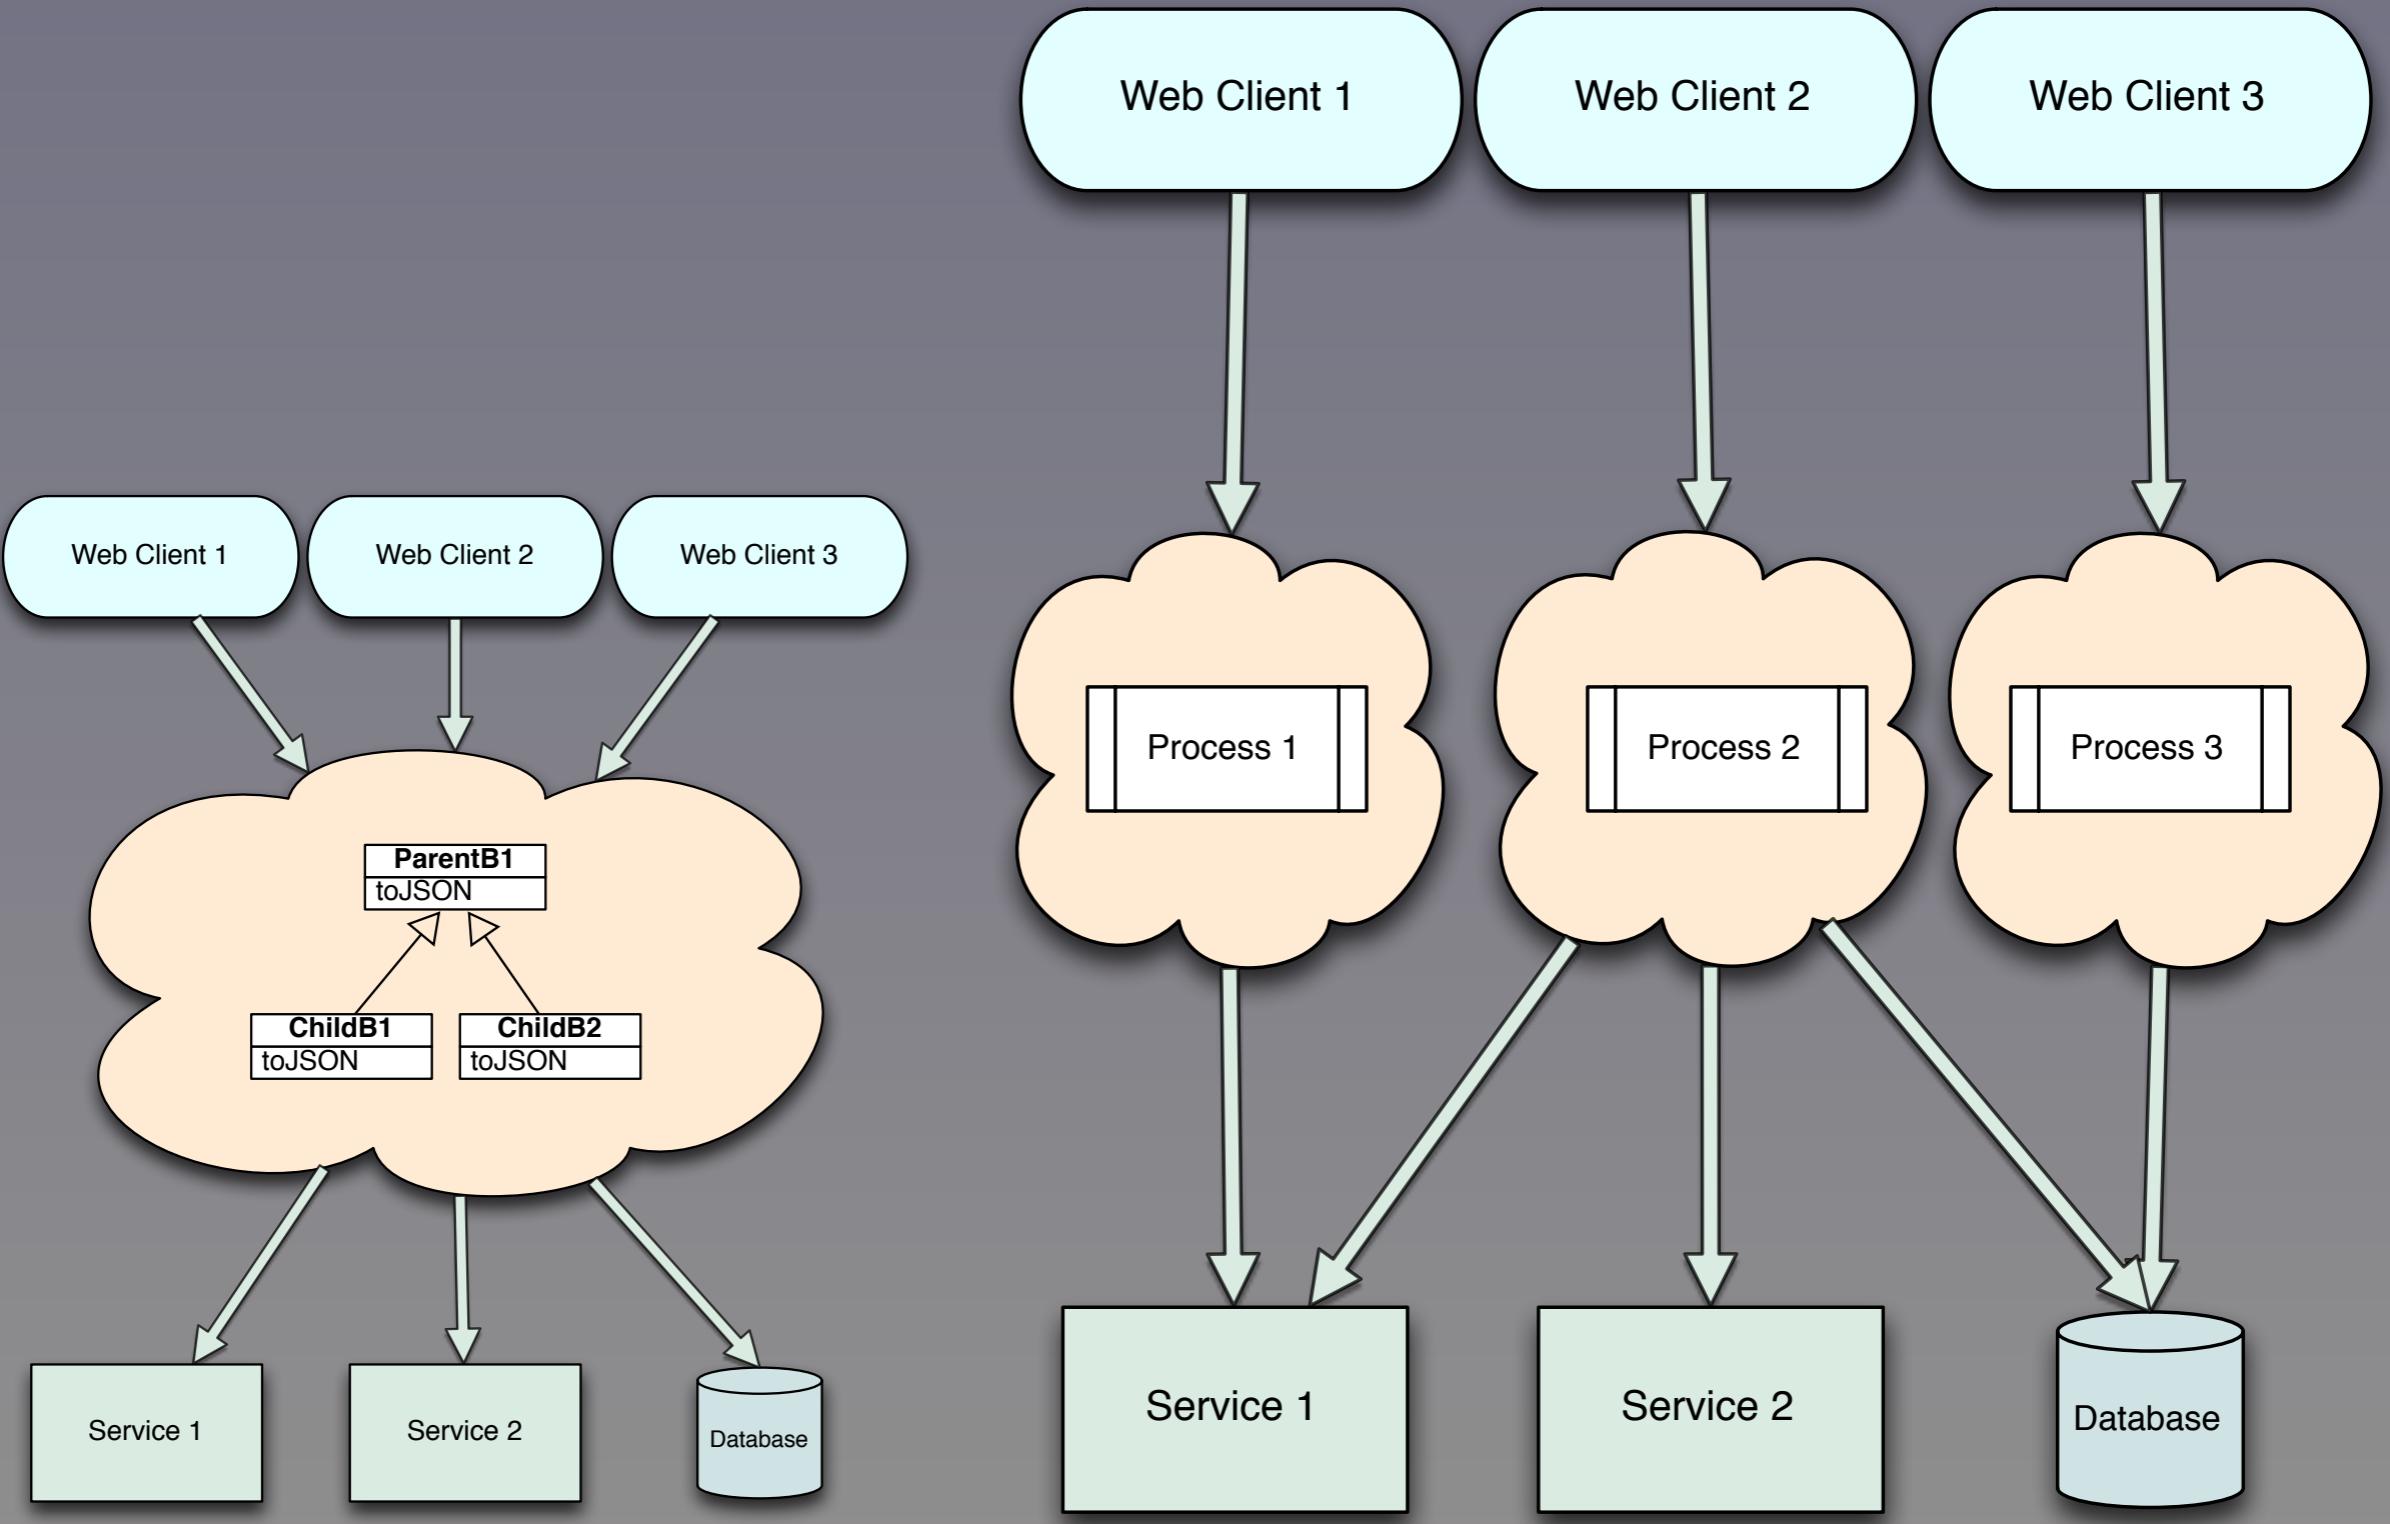
Tuesday, June 19, 12

This is the 2nd generation “Stanley”, the most successful self-driving car ever built (by a Google-Stanford) team. Machine learning is growing in importance. Here, generic algorithms and data structures are trained to represent the “world” using data, rather than encoding a model of the world in the software itself. It’s another example of generic algorithms that produce the desired behavior by being application agnostic and data driven, rather than hard-coding a model of the world. (In practice, however, a balance is struck between completely agnostic apps and some engineering towards the specific problem.)



Tuesday, June 19, 12

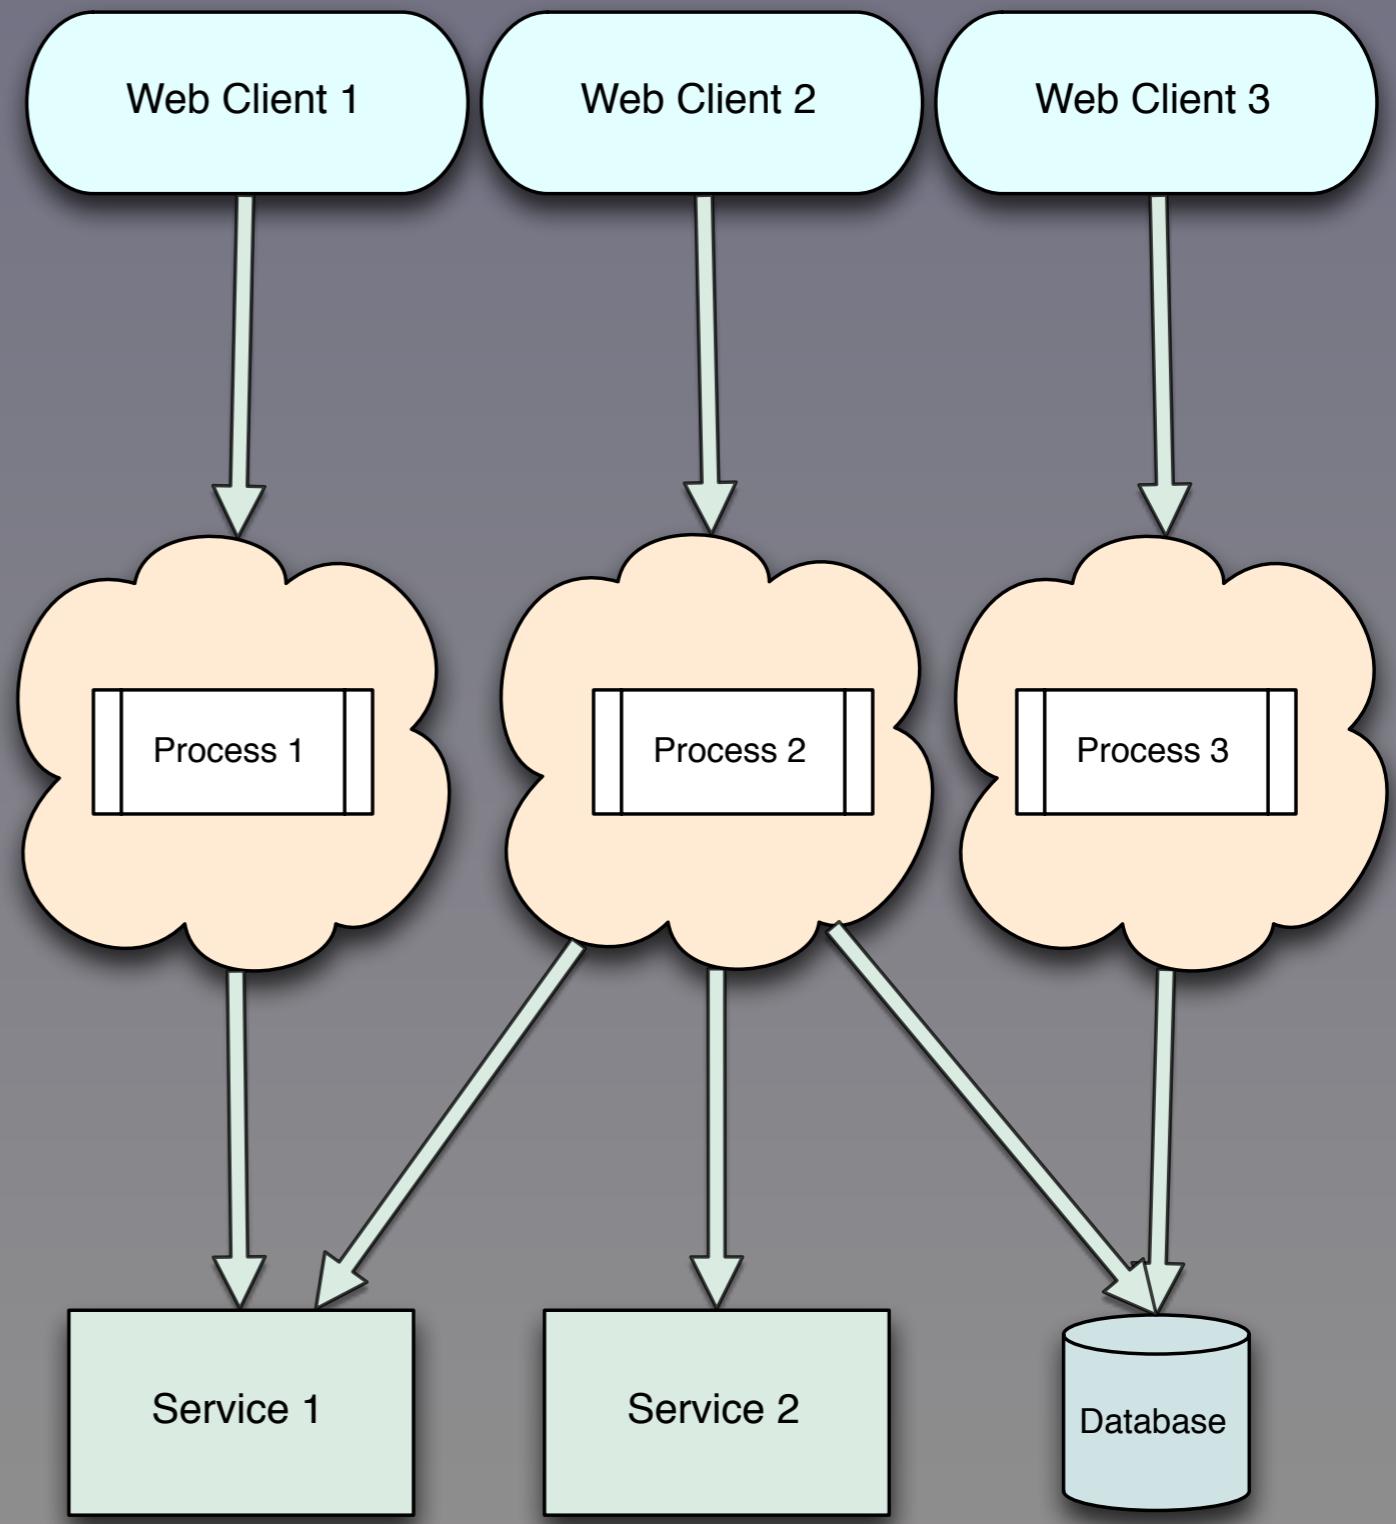
In a broader view, object models tend to push us towards centralized, complex systems that don't decompose well and stifle reuse and optimal deployment scenarios. FP code makes it easier to write smaller, focused services that we compose and deploy as appropriate.



Tuesday, June 19, 12

In a broader view, object models tend to push us towards centralized, complex systems that don't decompose well and stifle reuse and optimal deployment scenarios. FP code makes it easier to write smaller, focused services that we compose and deploy as appropriate.

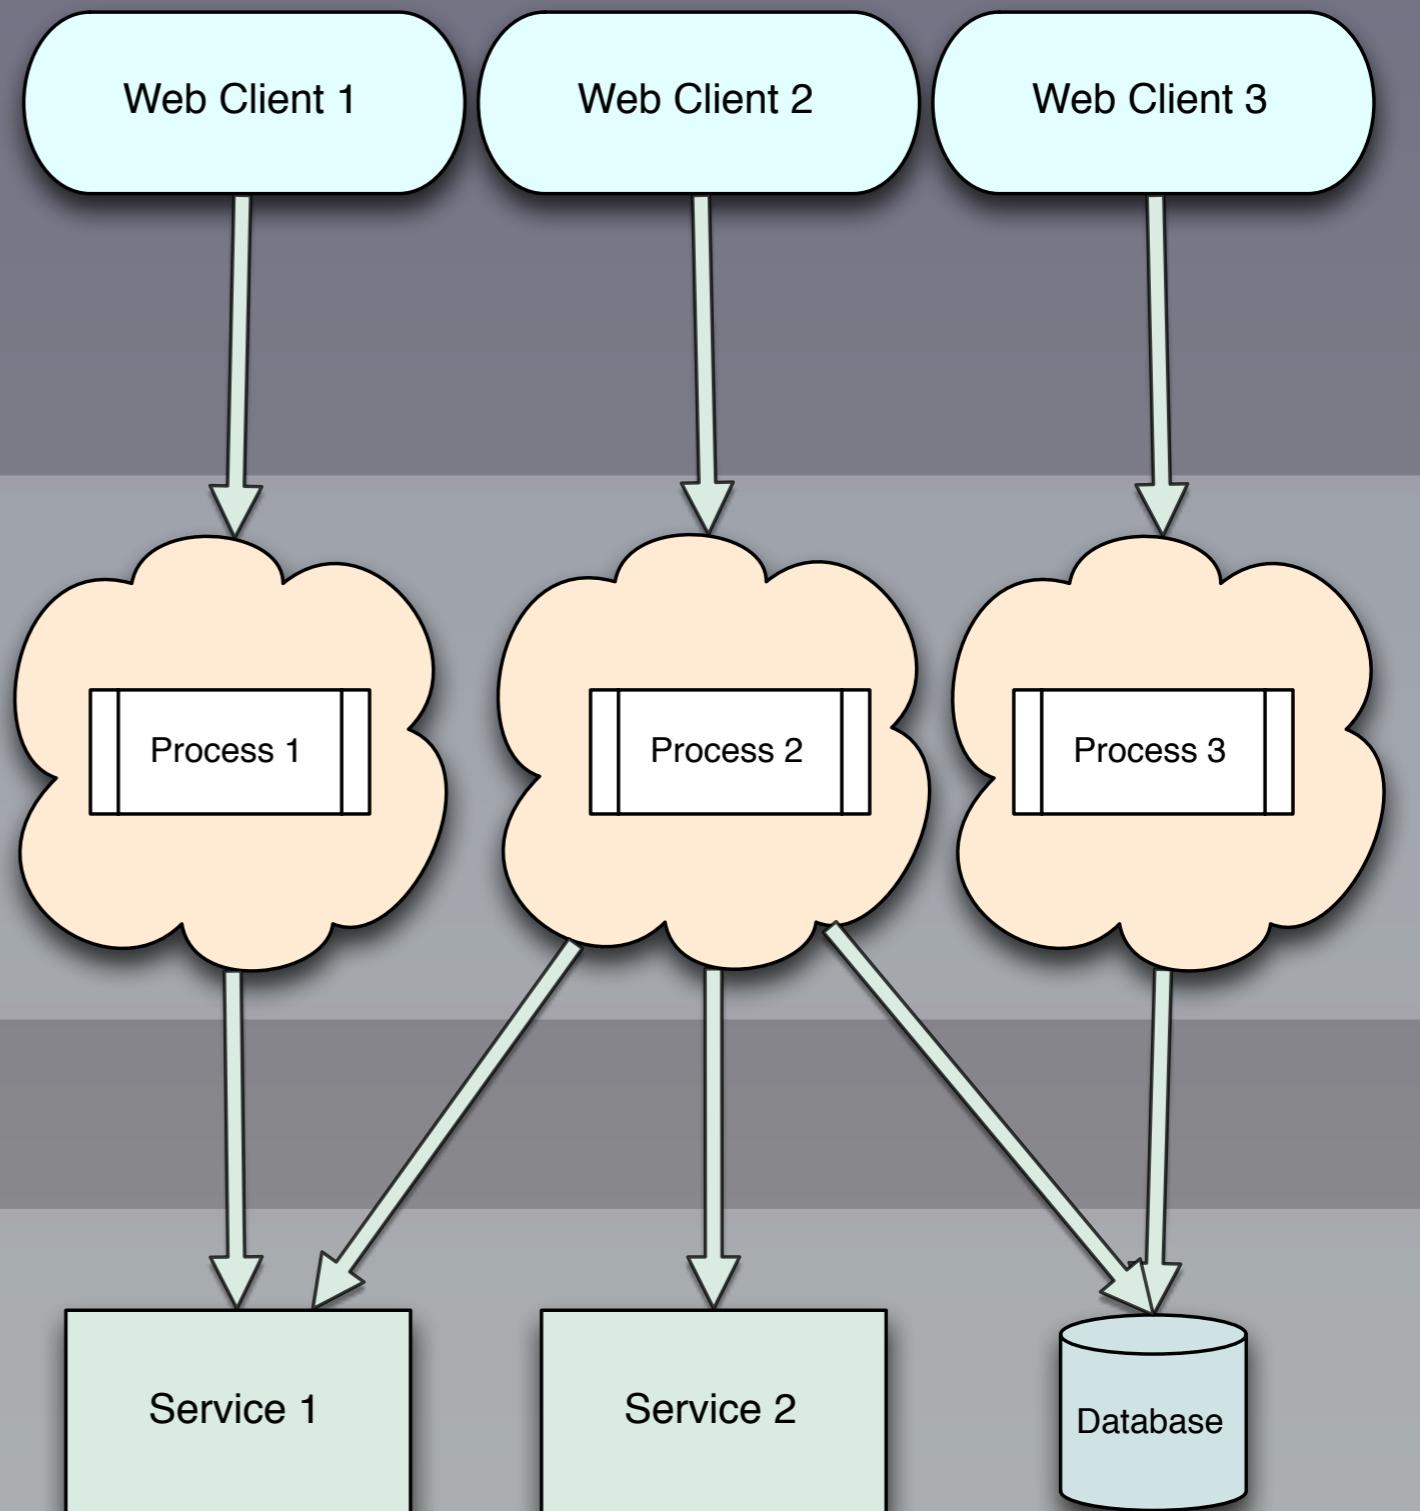
- Data Size ↑
- Formal Schema ↓
- Data-Driven Programs ↑



Tuesday, June 19, 12

Smaller, focused services scale better, especially horizontally. They also don't encapsulate more business logic than is required, and this (informal) architecture is also suitable for scaling ML and related algorithms.

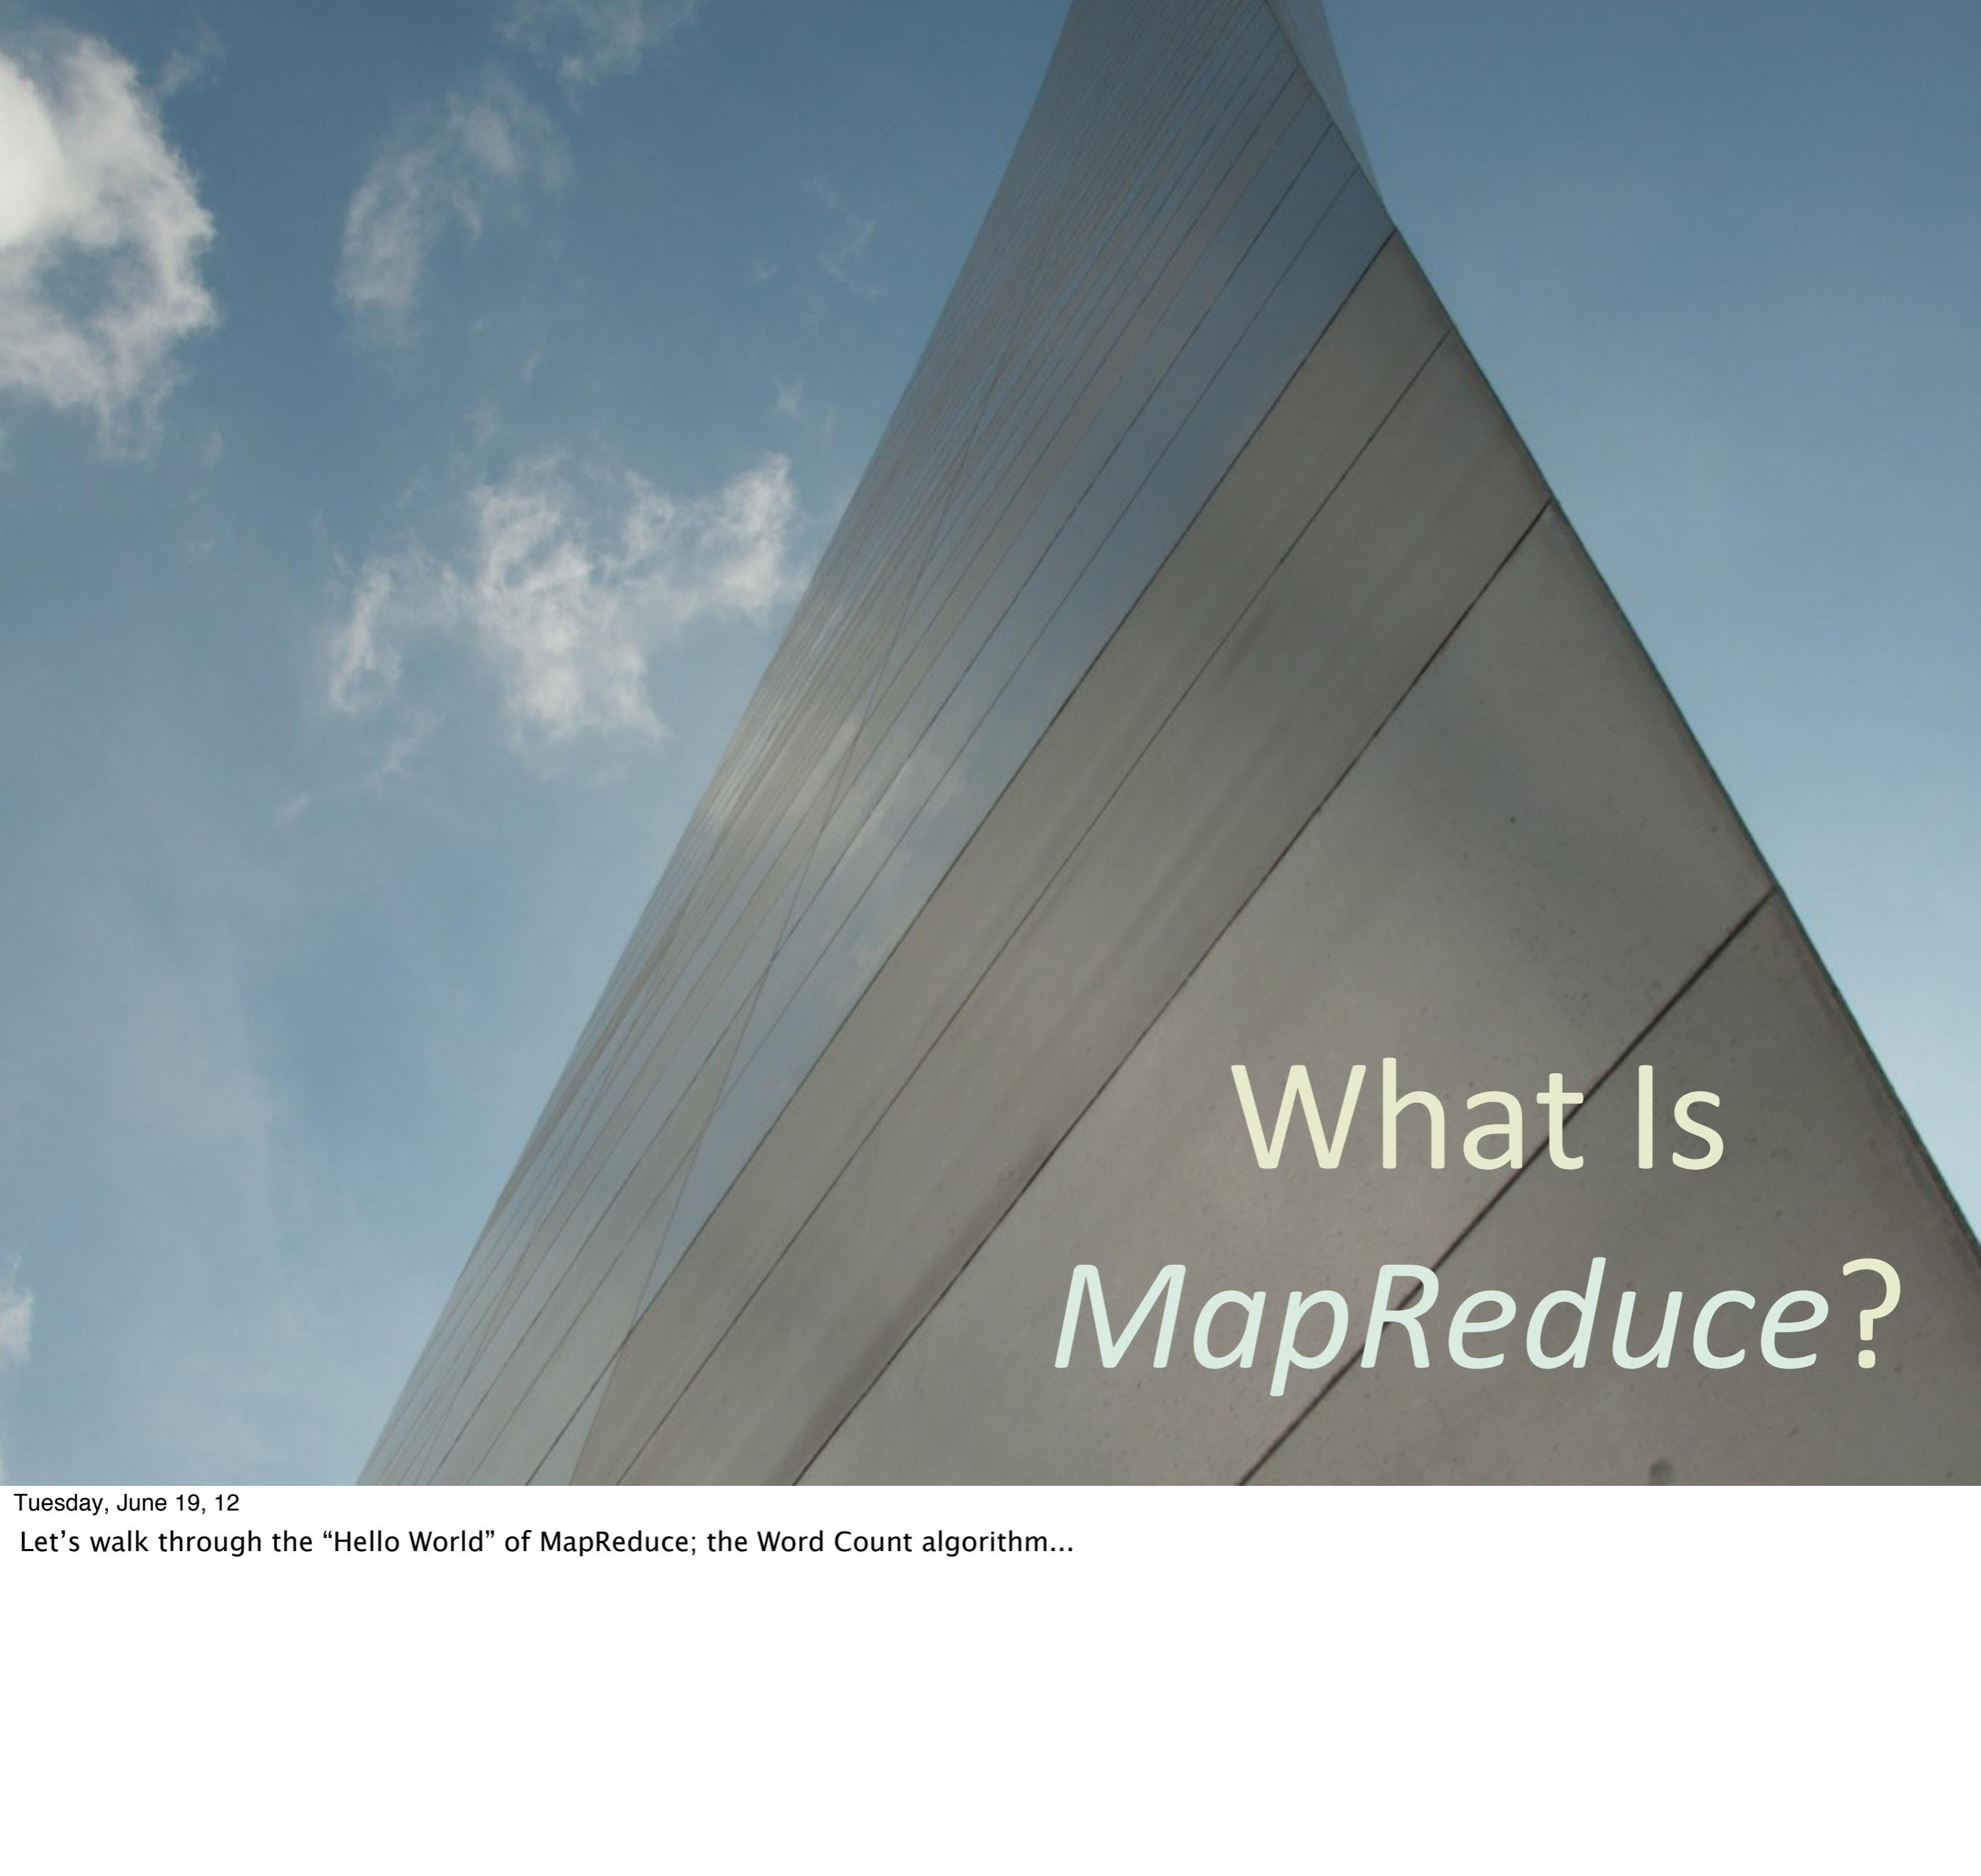
- MapReduce



- Distrib. FS

Tuesday, June 19, 12

And MapReduce + a distributed file system, like the Hadoop DFS, fit this model.

The background image shows a modern architectural structure with a large, curved, light-colored facade, possibly made of glass or a similar material. The building is set against a clear blue sky with some wispy white clouds. The perspective is from a low angle, looking up at the building's exterior.

What Is *MapReduce*?

Tuesday, June 19, 12

Let's walk through the "Hello World" of MapReduce; the Word Count algorithm...

MapReduce in Hadoop

Let's look at a
MapReduce algorithm:
WordCount.

(The *Hello World* of big data...)

Tuesday, June 19, 12

Let's walk through this algorithm at a conceptual level...

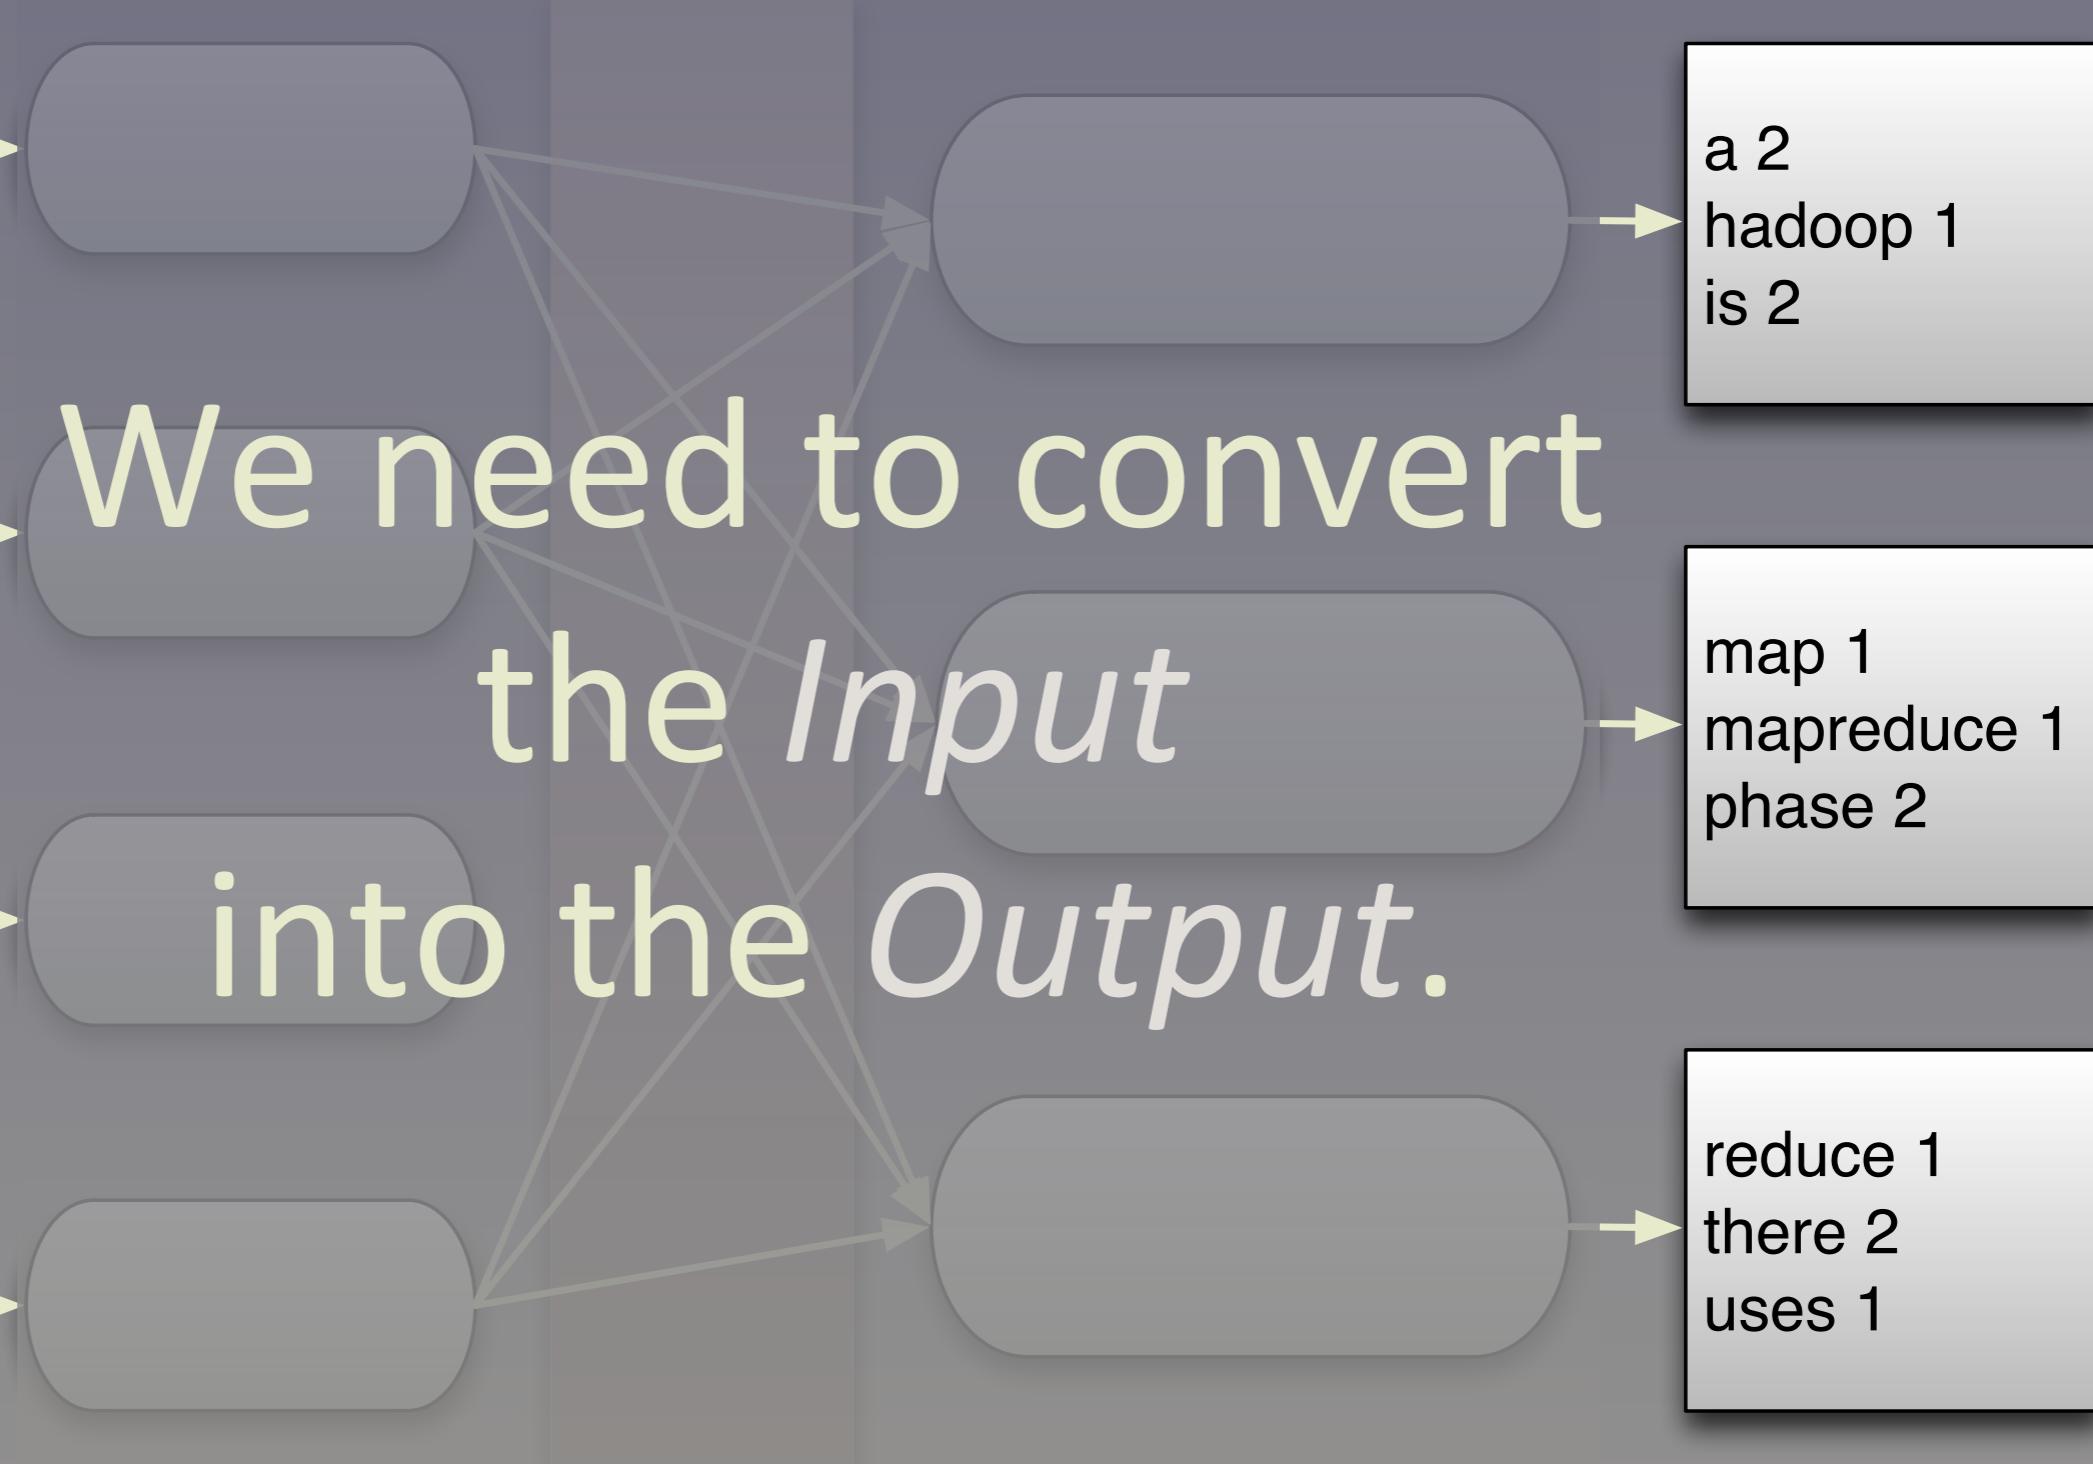
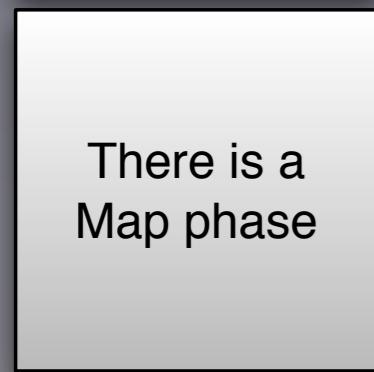
Input

Mappers

Sort,
Shuffle

Reducers

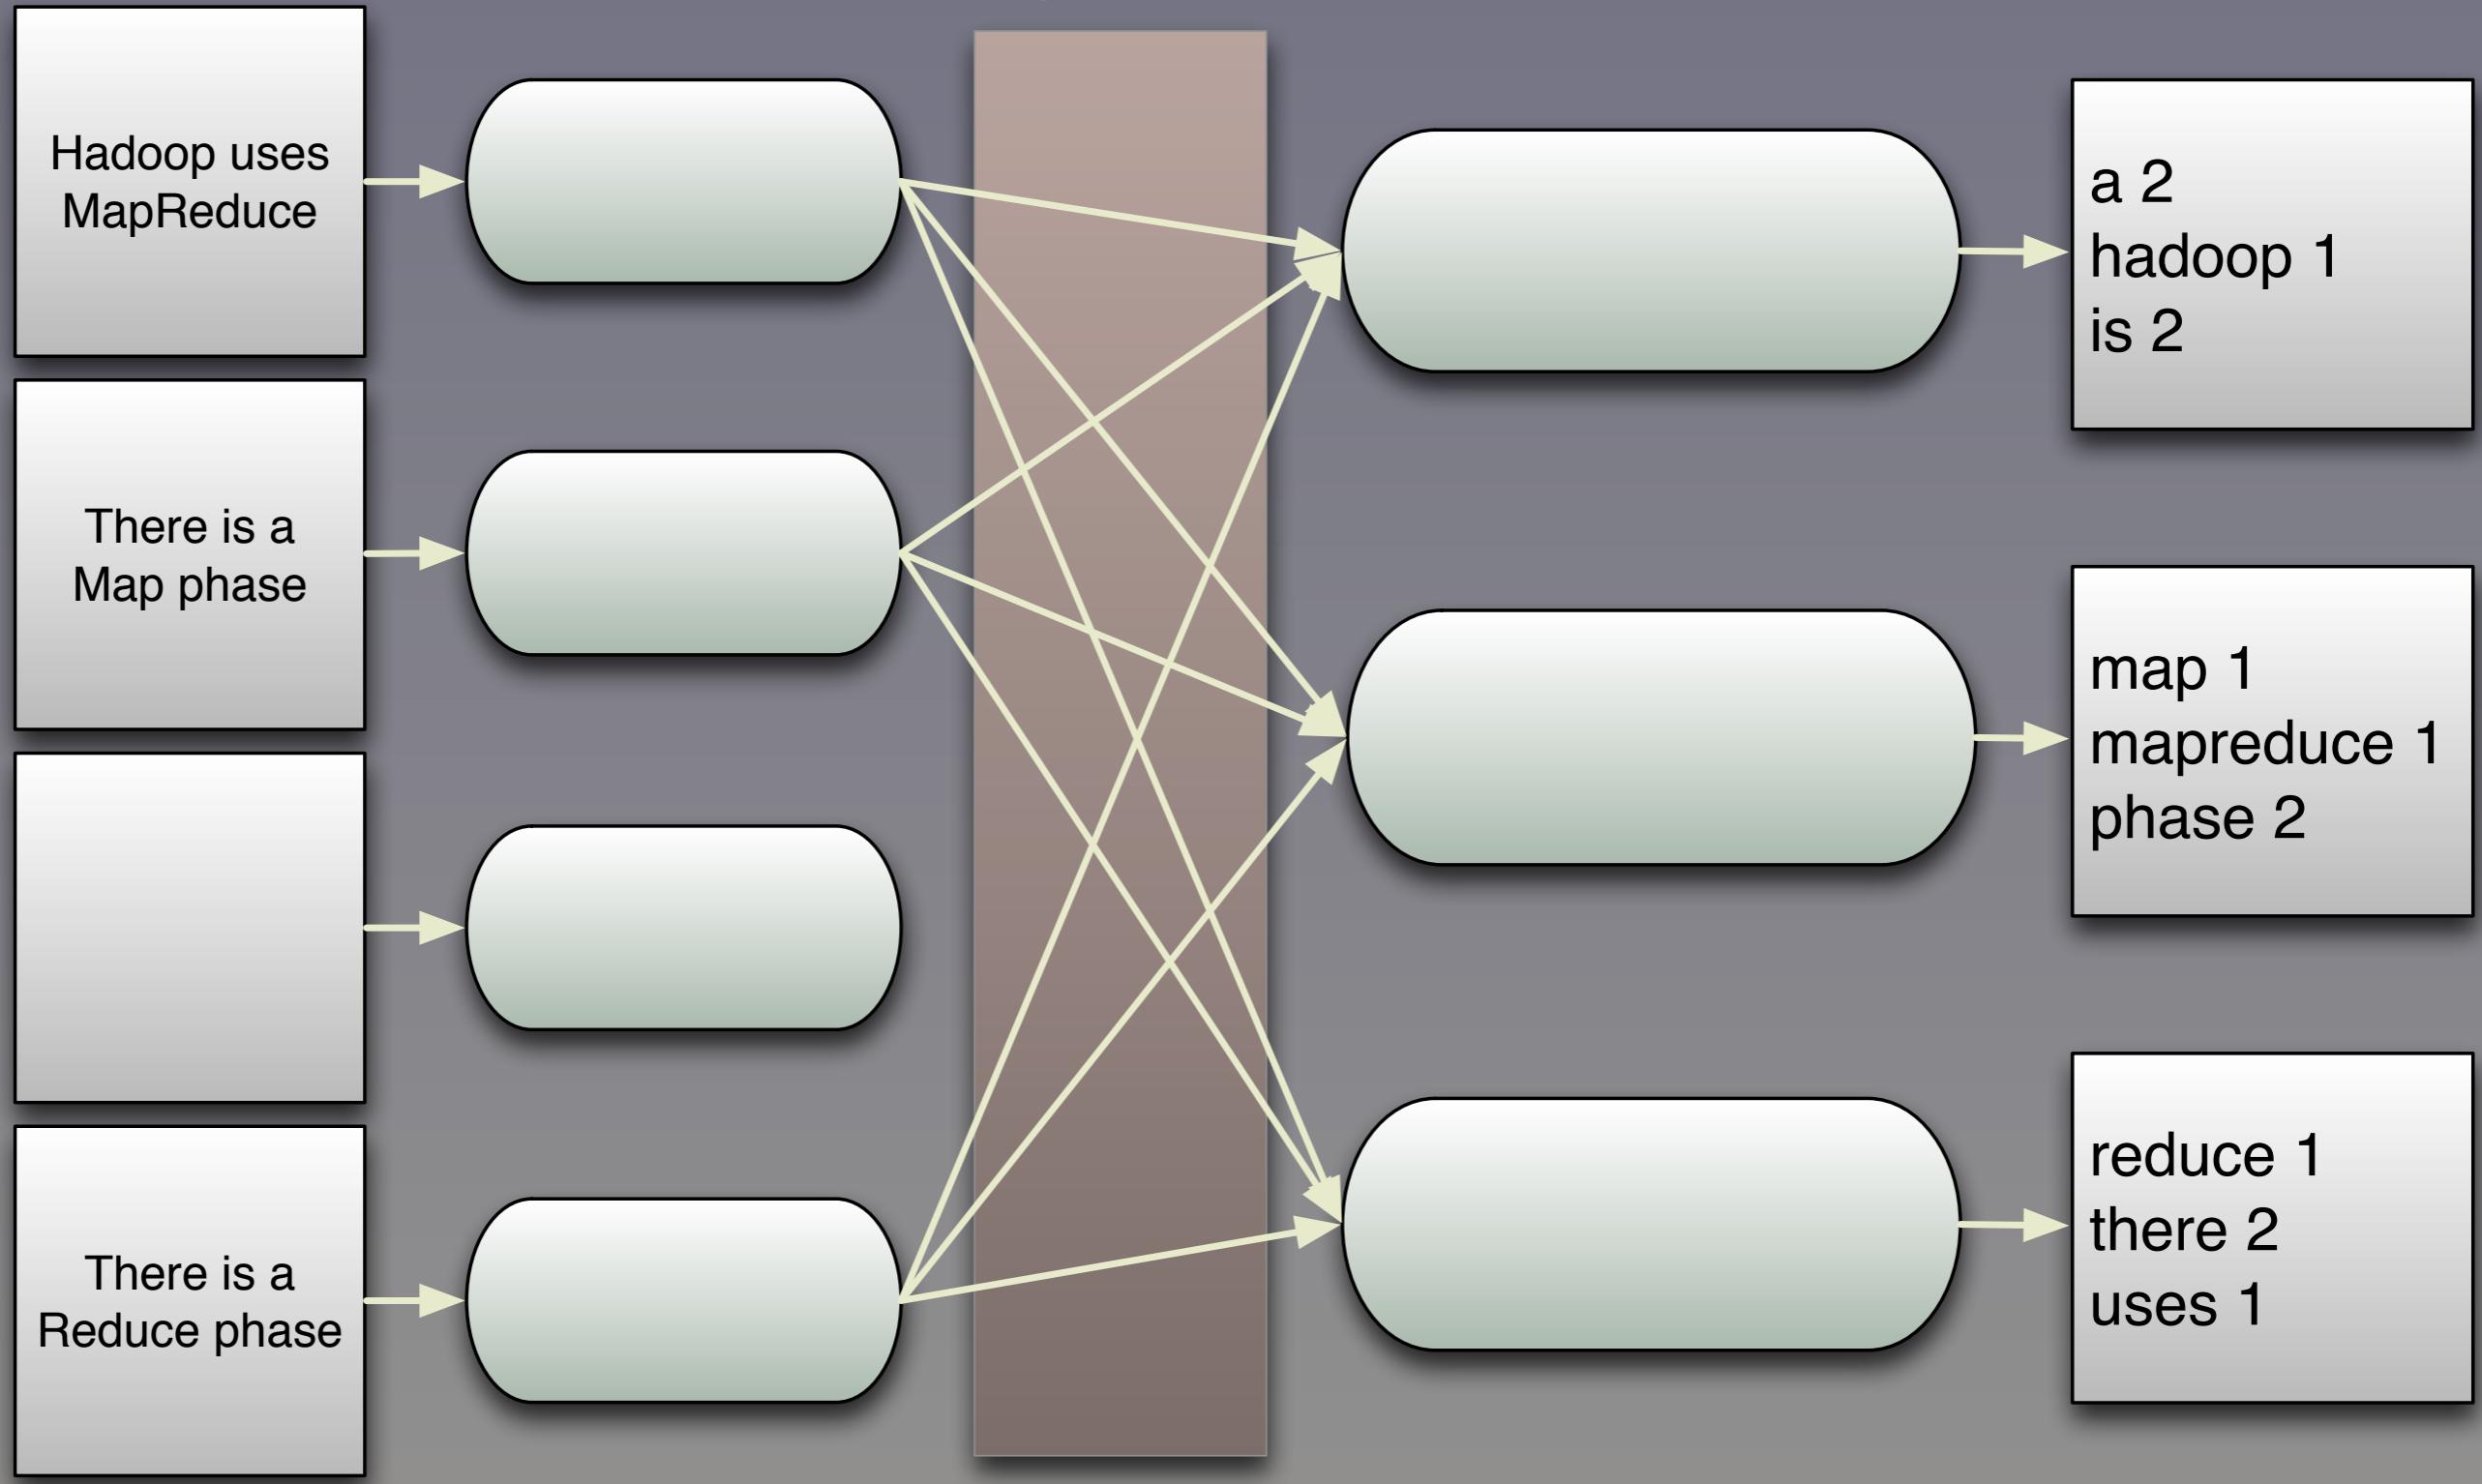
Output



Tuesday, June 19, 12

Four input documents, one left empty, the others with small phrases (for simplicity...). The word count output is on the right (we'll see why there are three output "documents"). We need to get from the input on the left-hand side to the output on the right-hand side.

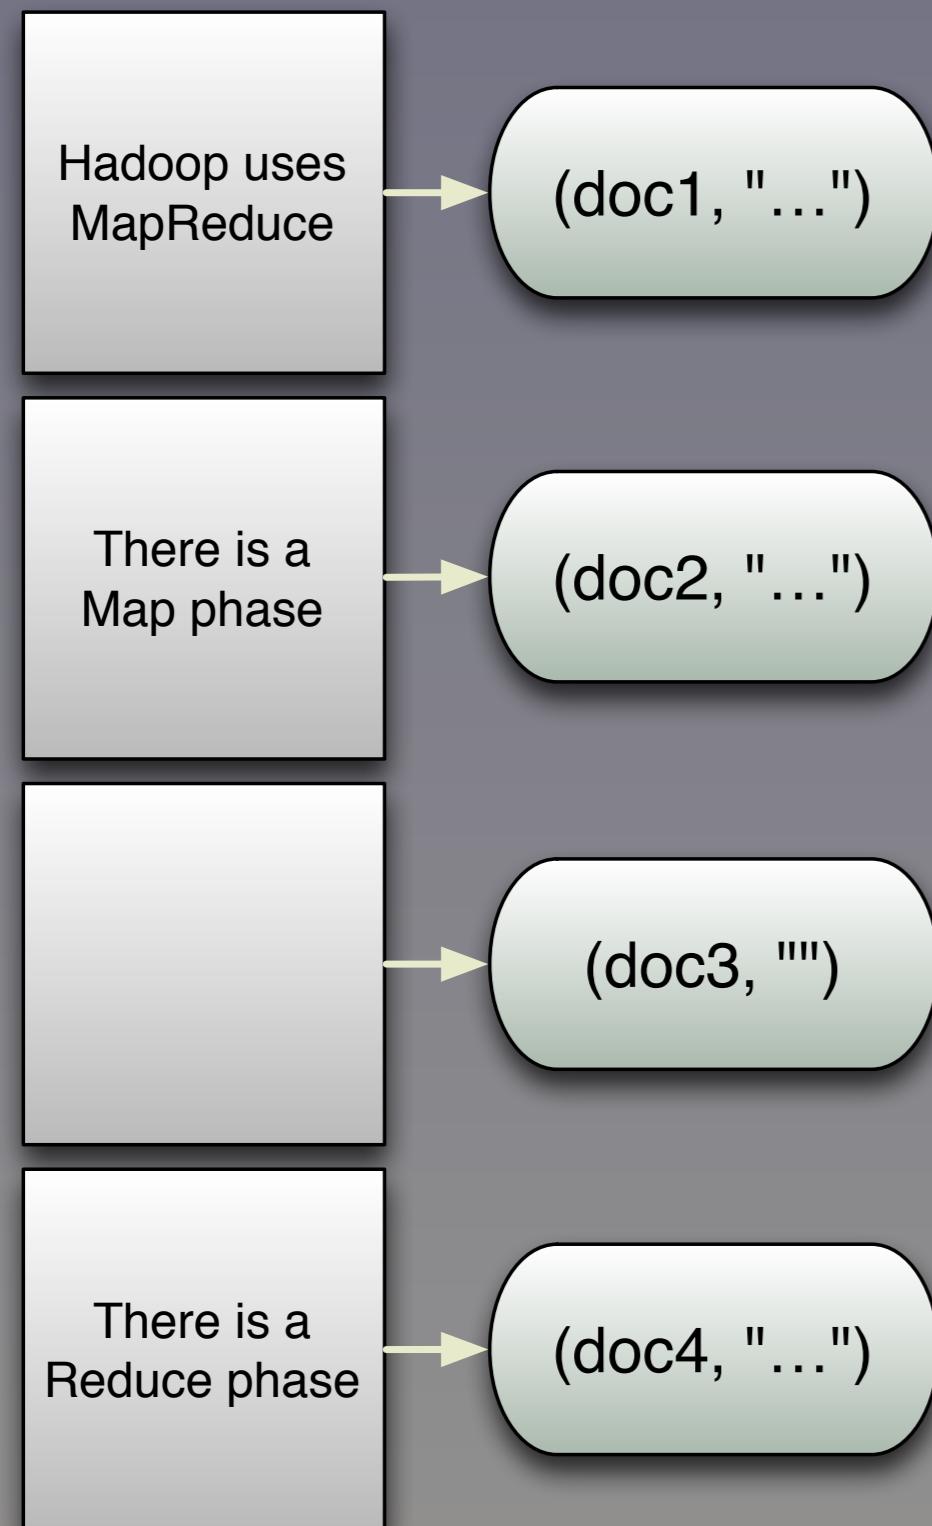
Input Mappers Sort, Shuffle Reducers Output



Tuesday, June 19, 12

Here is a schematic view of the steps in Hadoop MapReduce. Each Input file is read by a single Mapper process (default: can be many-to-many, as we'll see later). The Mappers emit key-value pairs that will be sorted, then partitioned and “shuffled” to the reducers, where each Reducer will get all instances of a given key (for 1 or more values). Each Reducer generates the final key-value pairs and writes them to one or more files (based on the size of the output).

Input Mappers

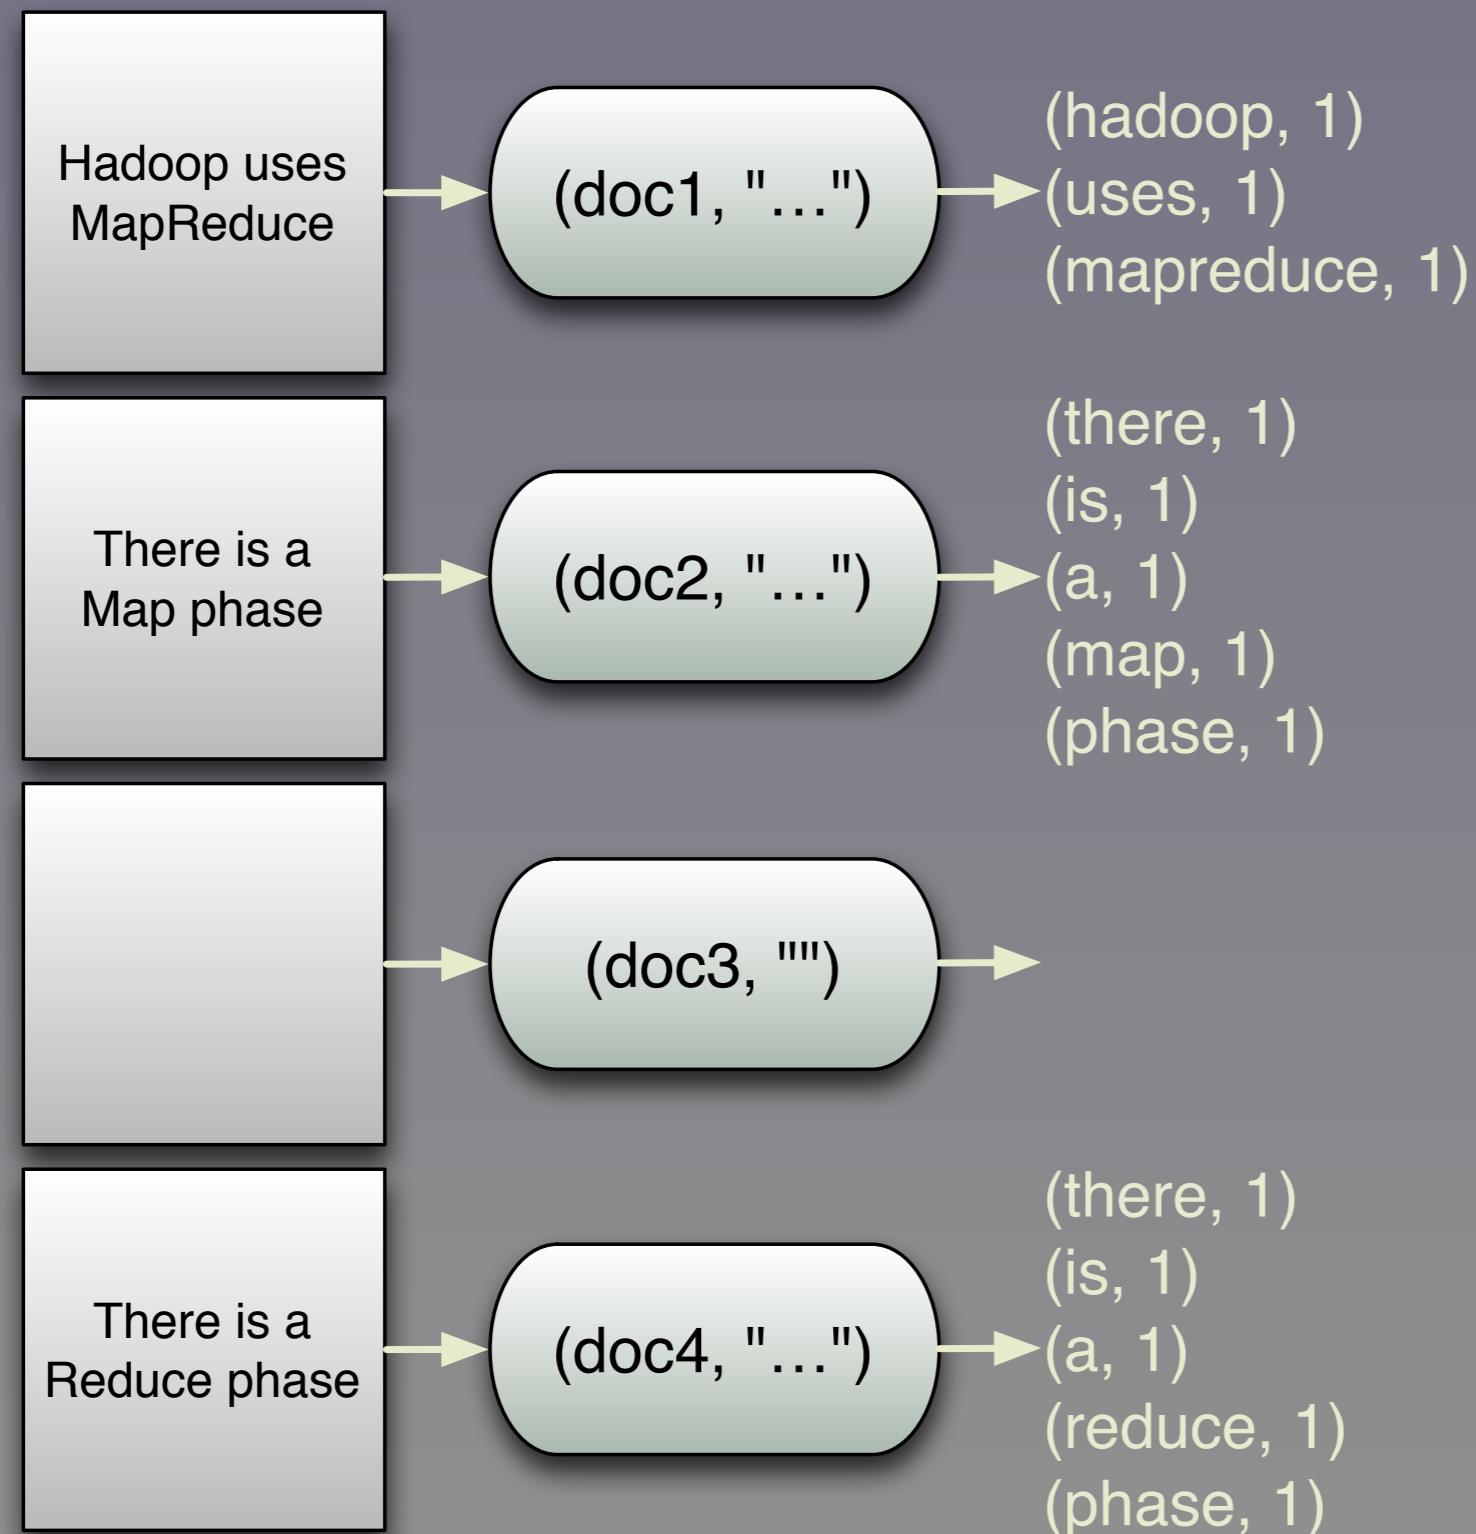


Tuesday, June 19, 12

Each document gets a mapper. All data is organized into key-value pairs. We don't care about the key. The value is all the content or line by line. (Doesn't matter.) I'm showing the document contents in the boxes for this example. Actually, large documents might get split to several mappers (as we'll see). It is also possible to concatenate many small documents into a single, larger document for input to a mapper.

Each mapper will tokenize the doc contents, convert all words to lower case and count them...

Input Mappers



Tuesday, June 19, 12

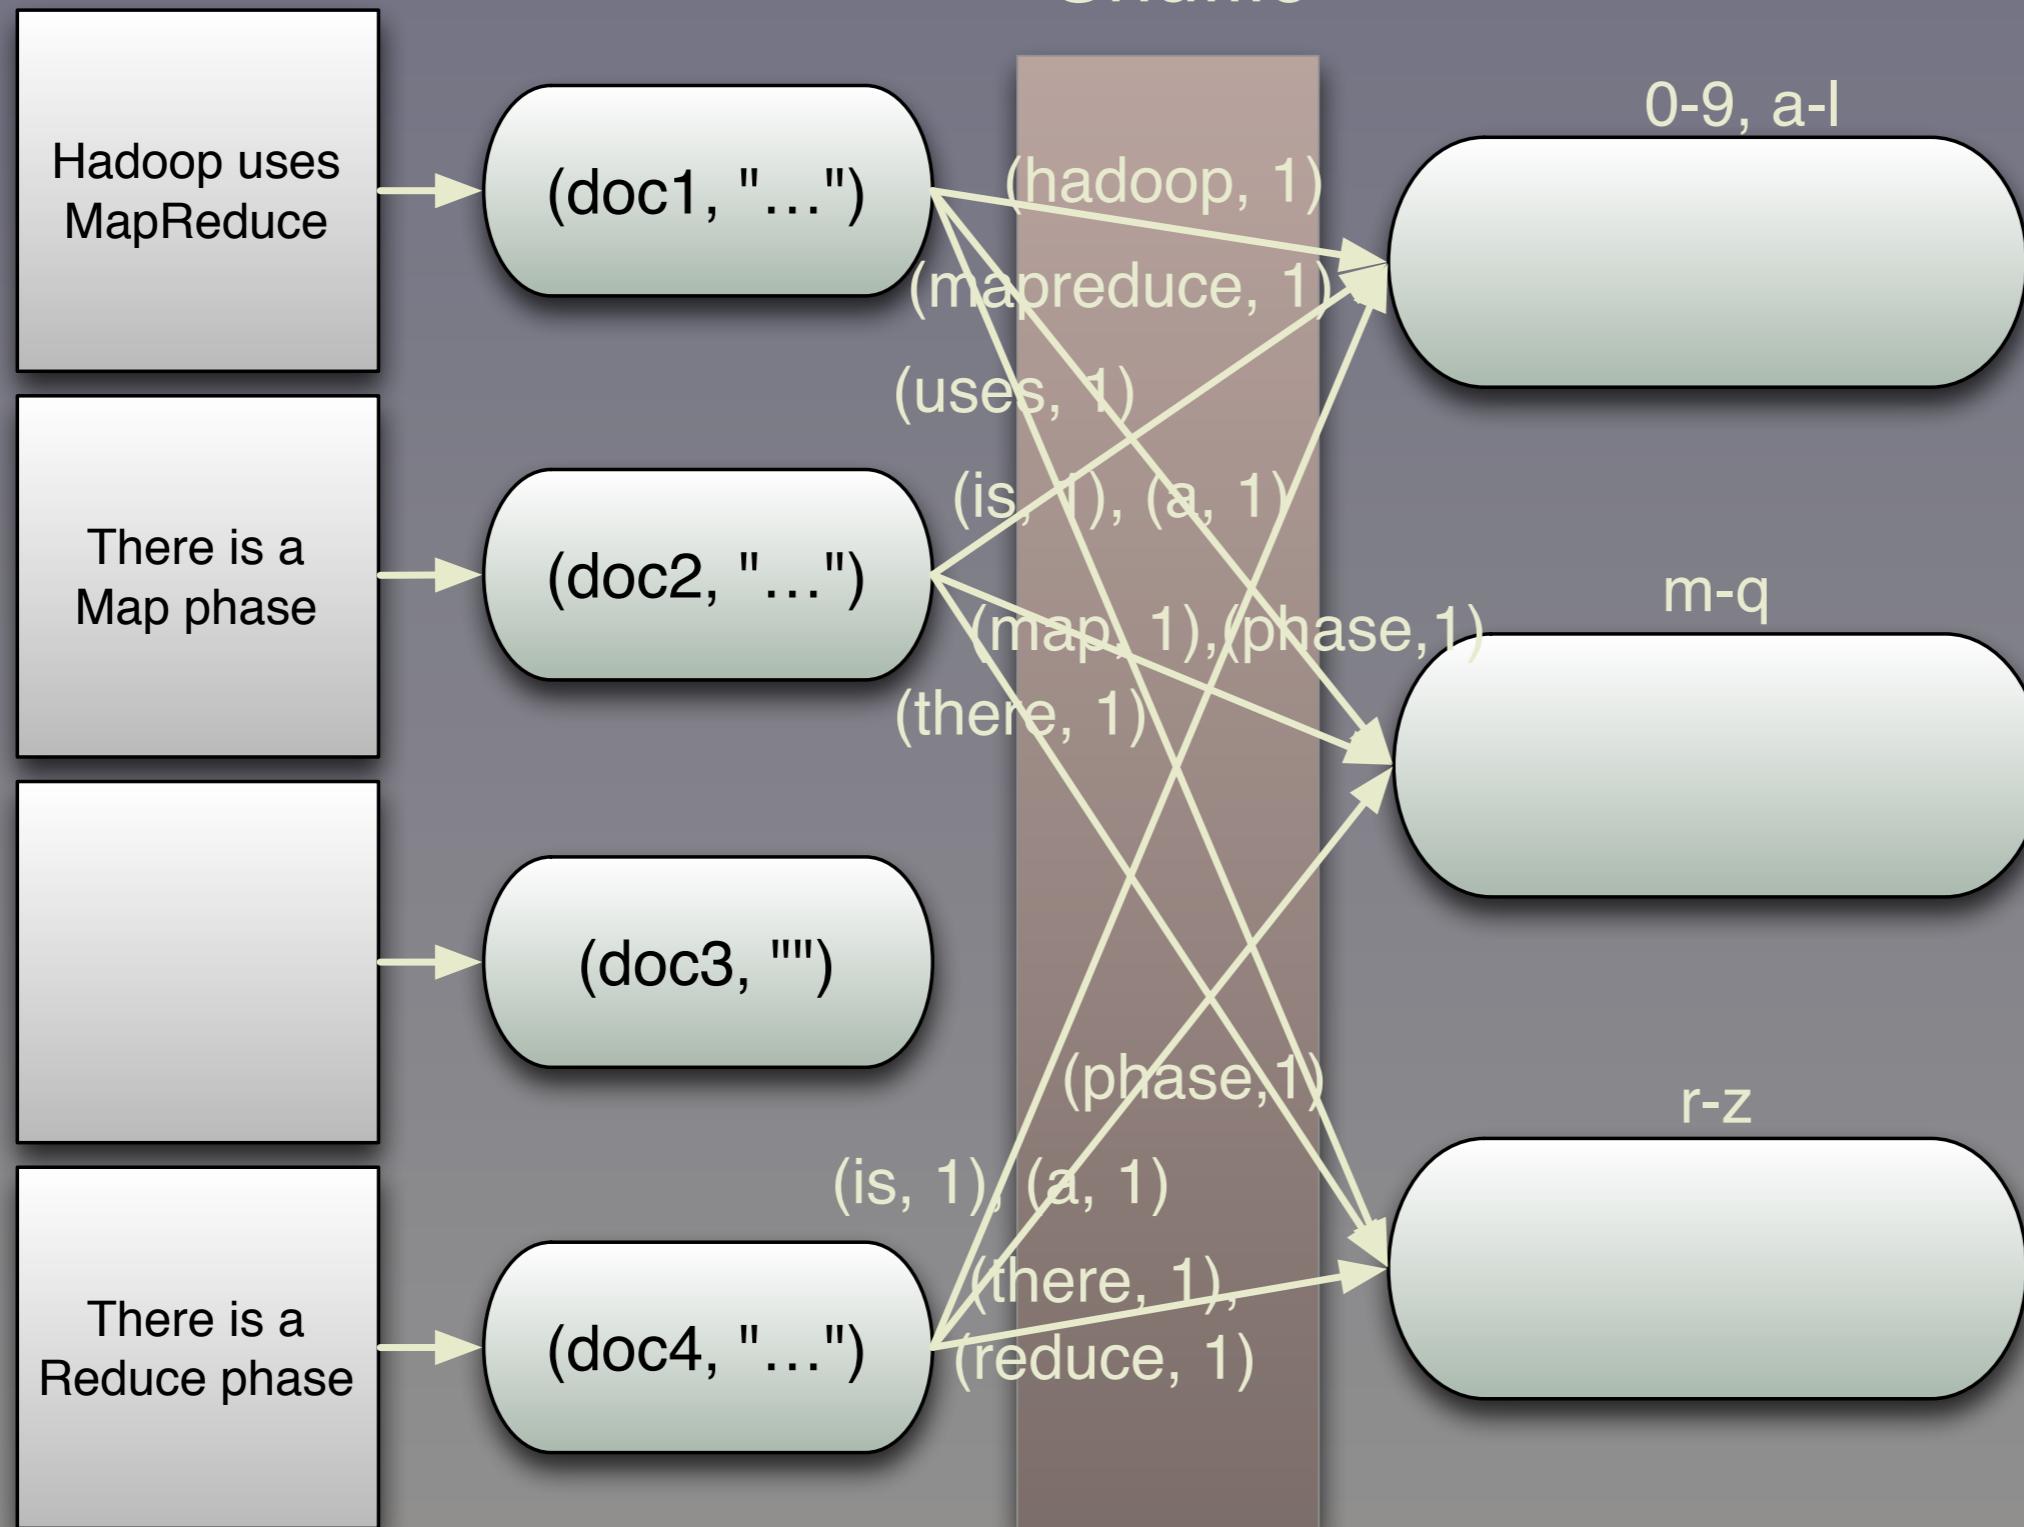
The mappers emit key-value pairs, where each key is one of the words, and the value is the count. In the most naive (but also most memory efficient) implementation, each mapper simply emits (word, 1) each time “word” is seen. However, this is IO inefficient! Note that the mapper for the empty doc. emits no pairs, as you would expect.

Input

Mappers

Sort, Shuffle

Reducers

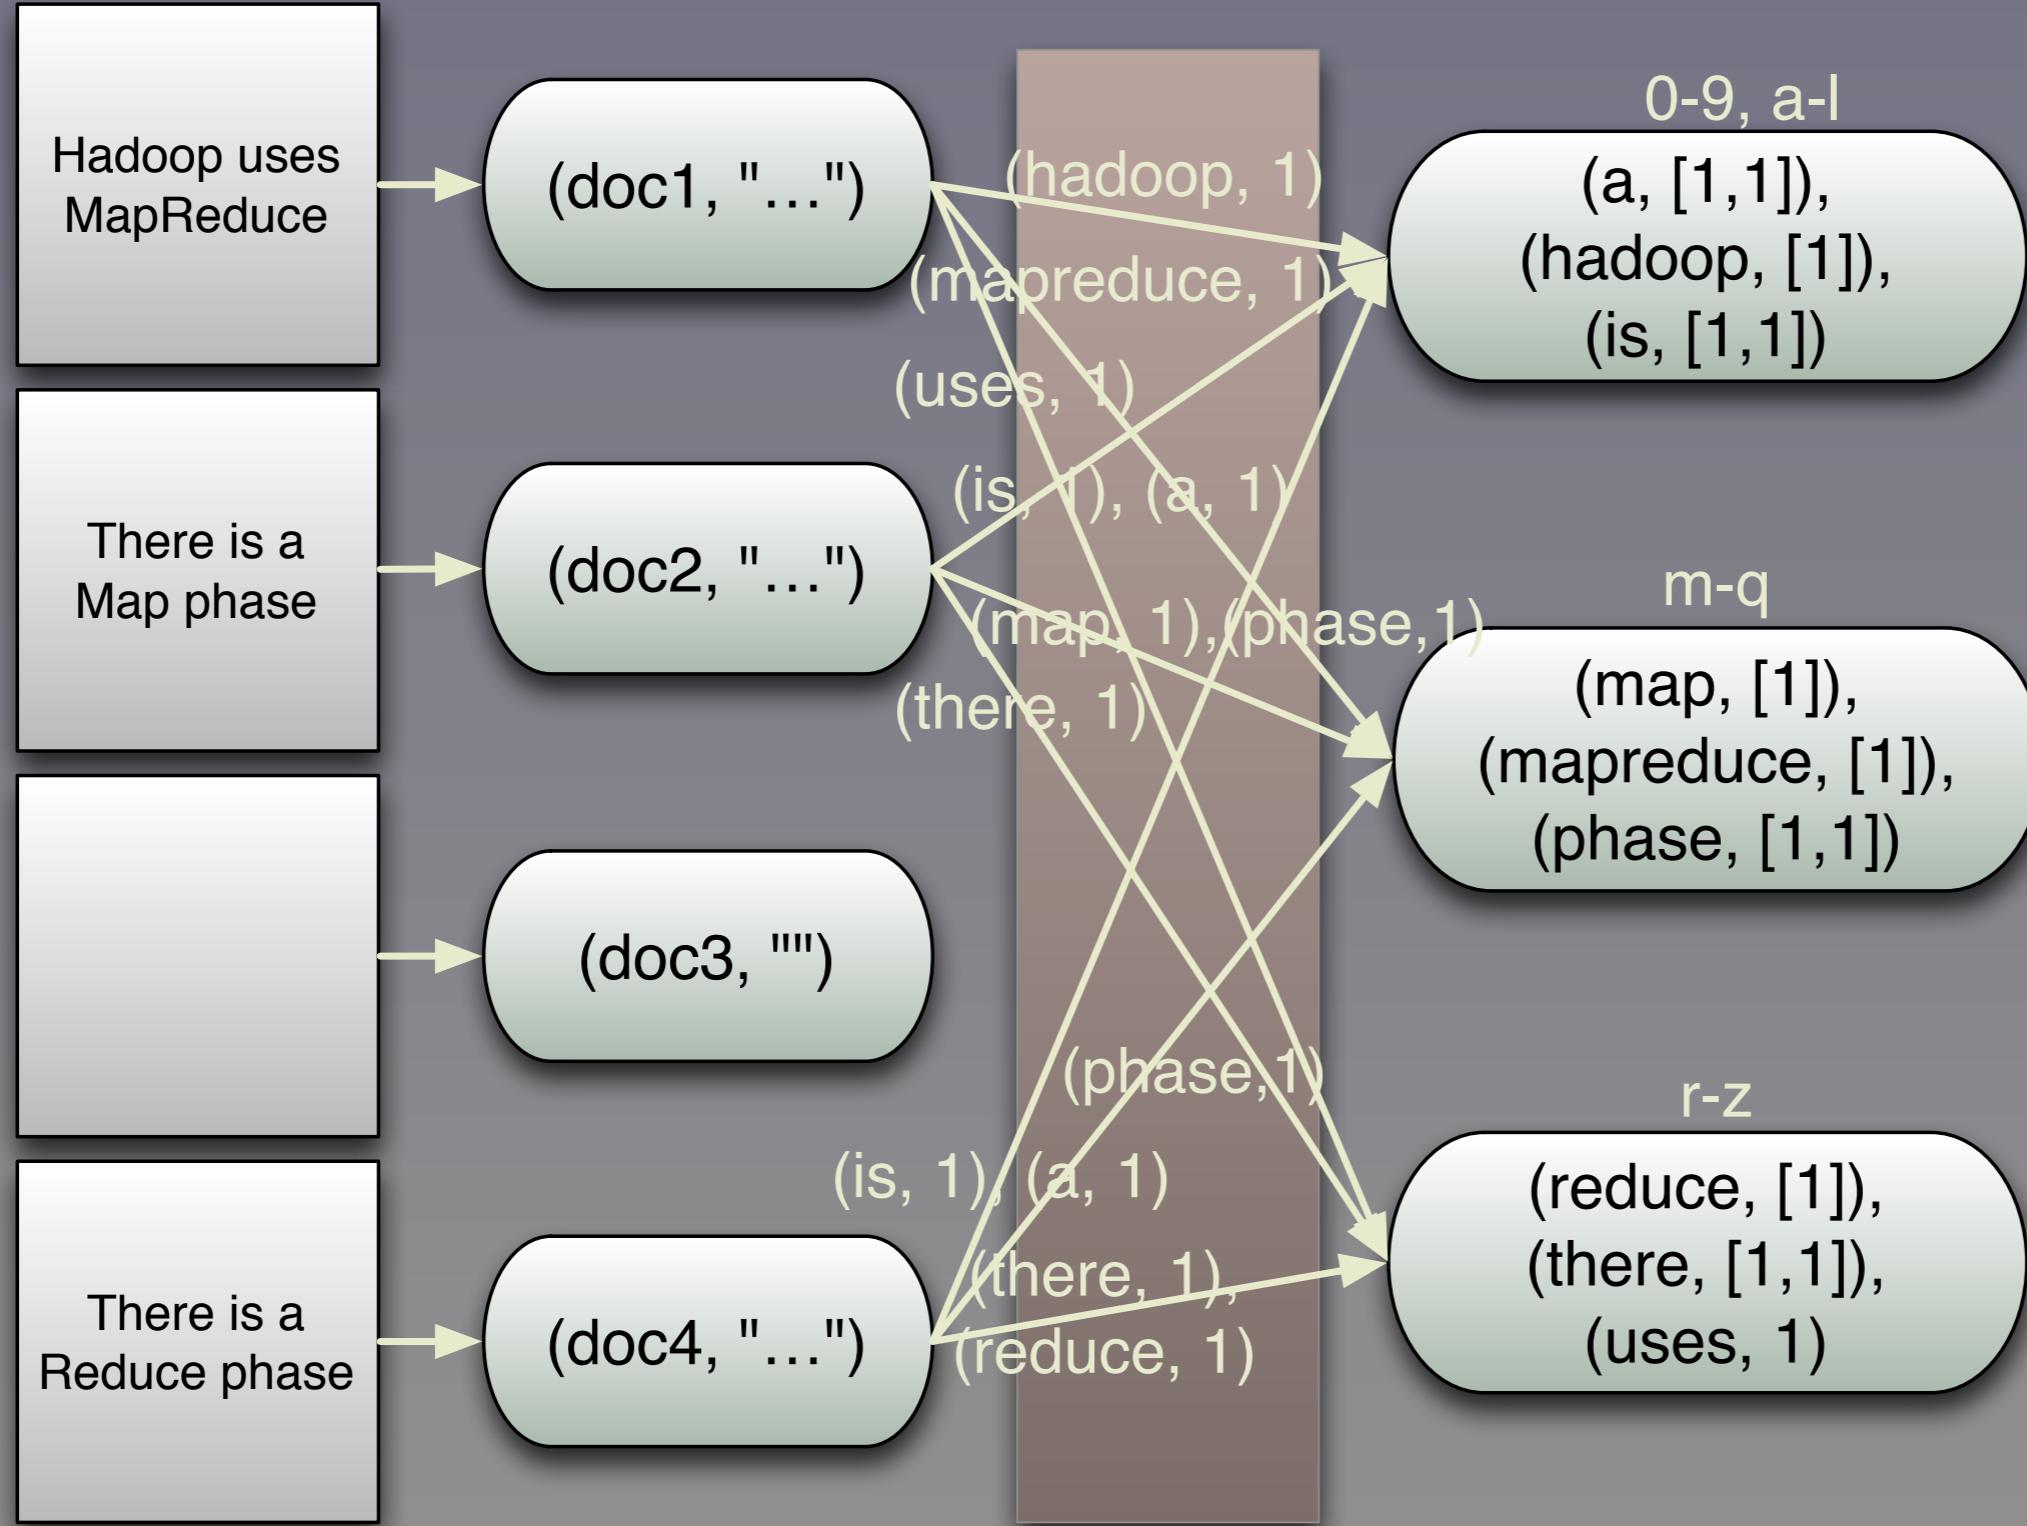


Tuesday, June 19, 12

The mappers themselves don't decide to which reducer each pair should be sent. Rather, the job setup configures what to do and the Hadoop runtime enforces it during the Sort/Shuffle phase, where the key-value pairs in each mapper are sorted by key (that is locally, not globally) and then the pairs are routed to the correct reducer, on the current machine or other machines.

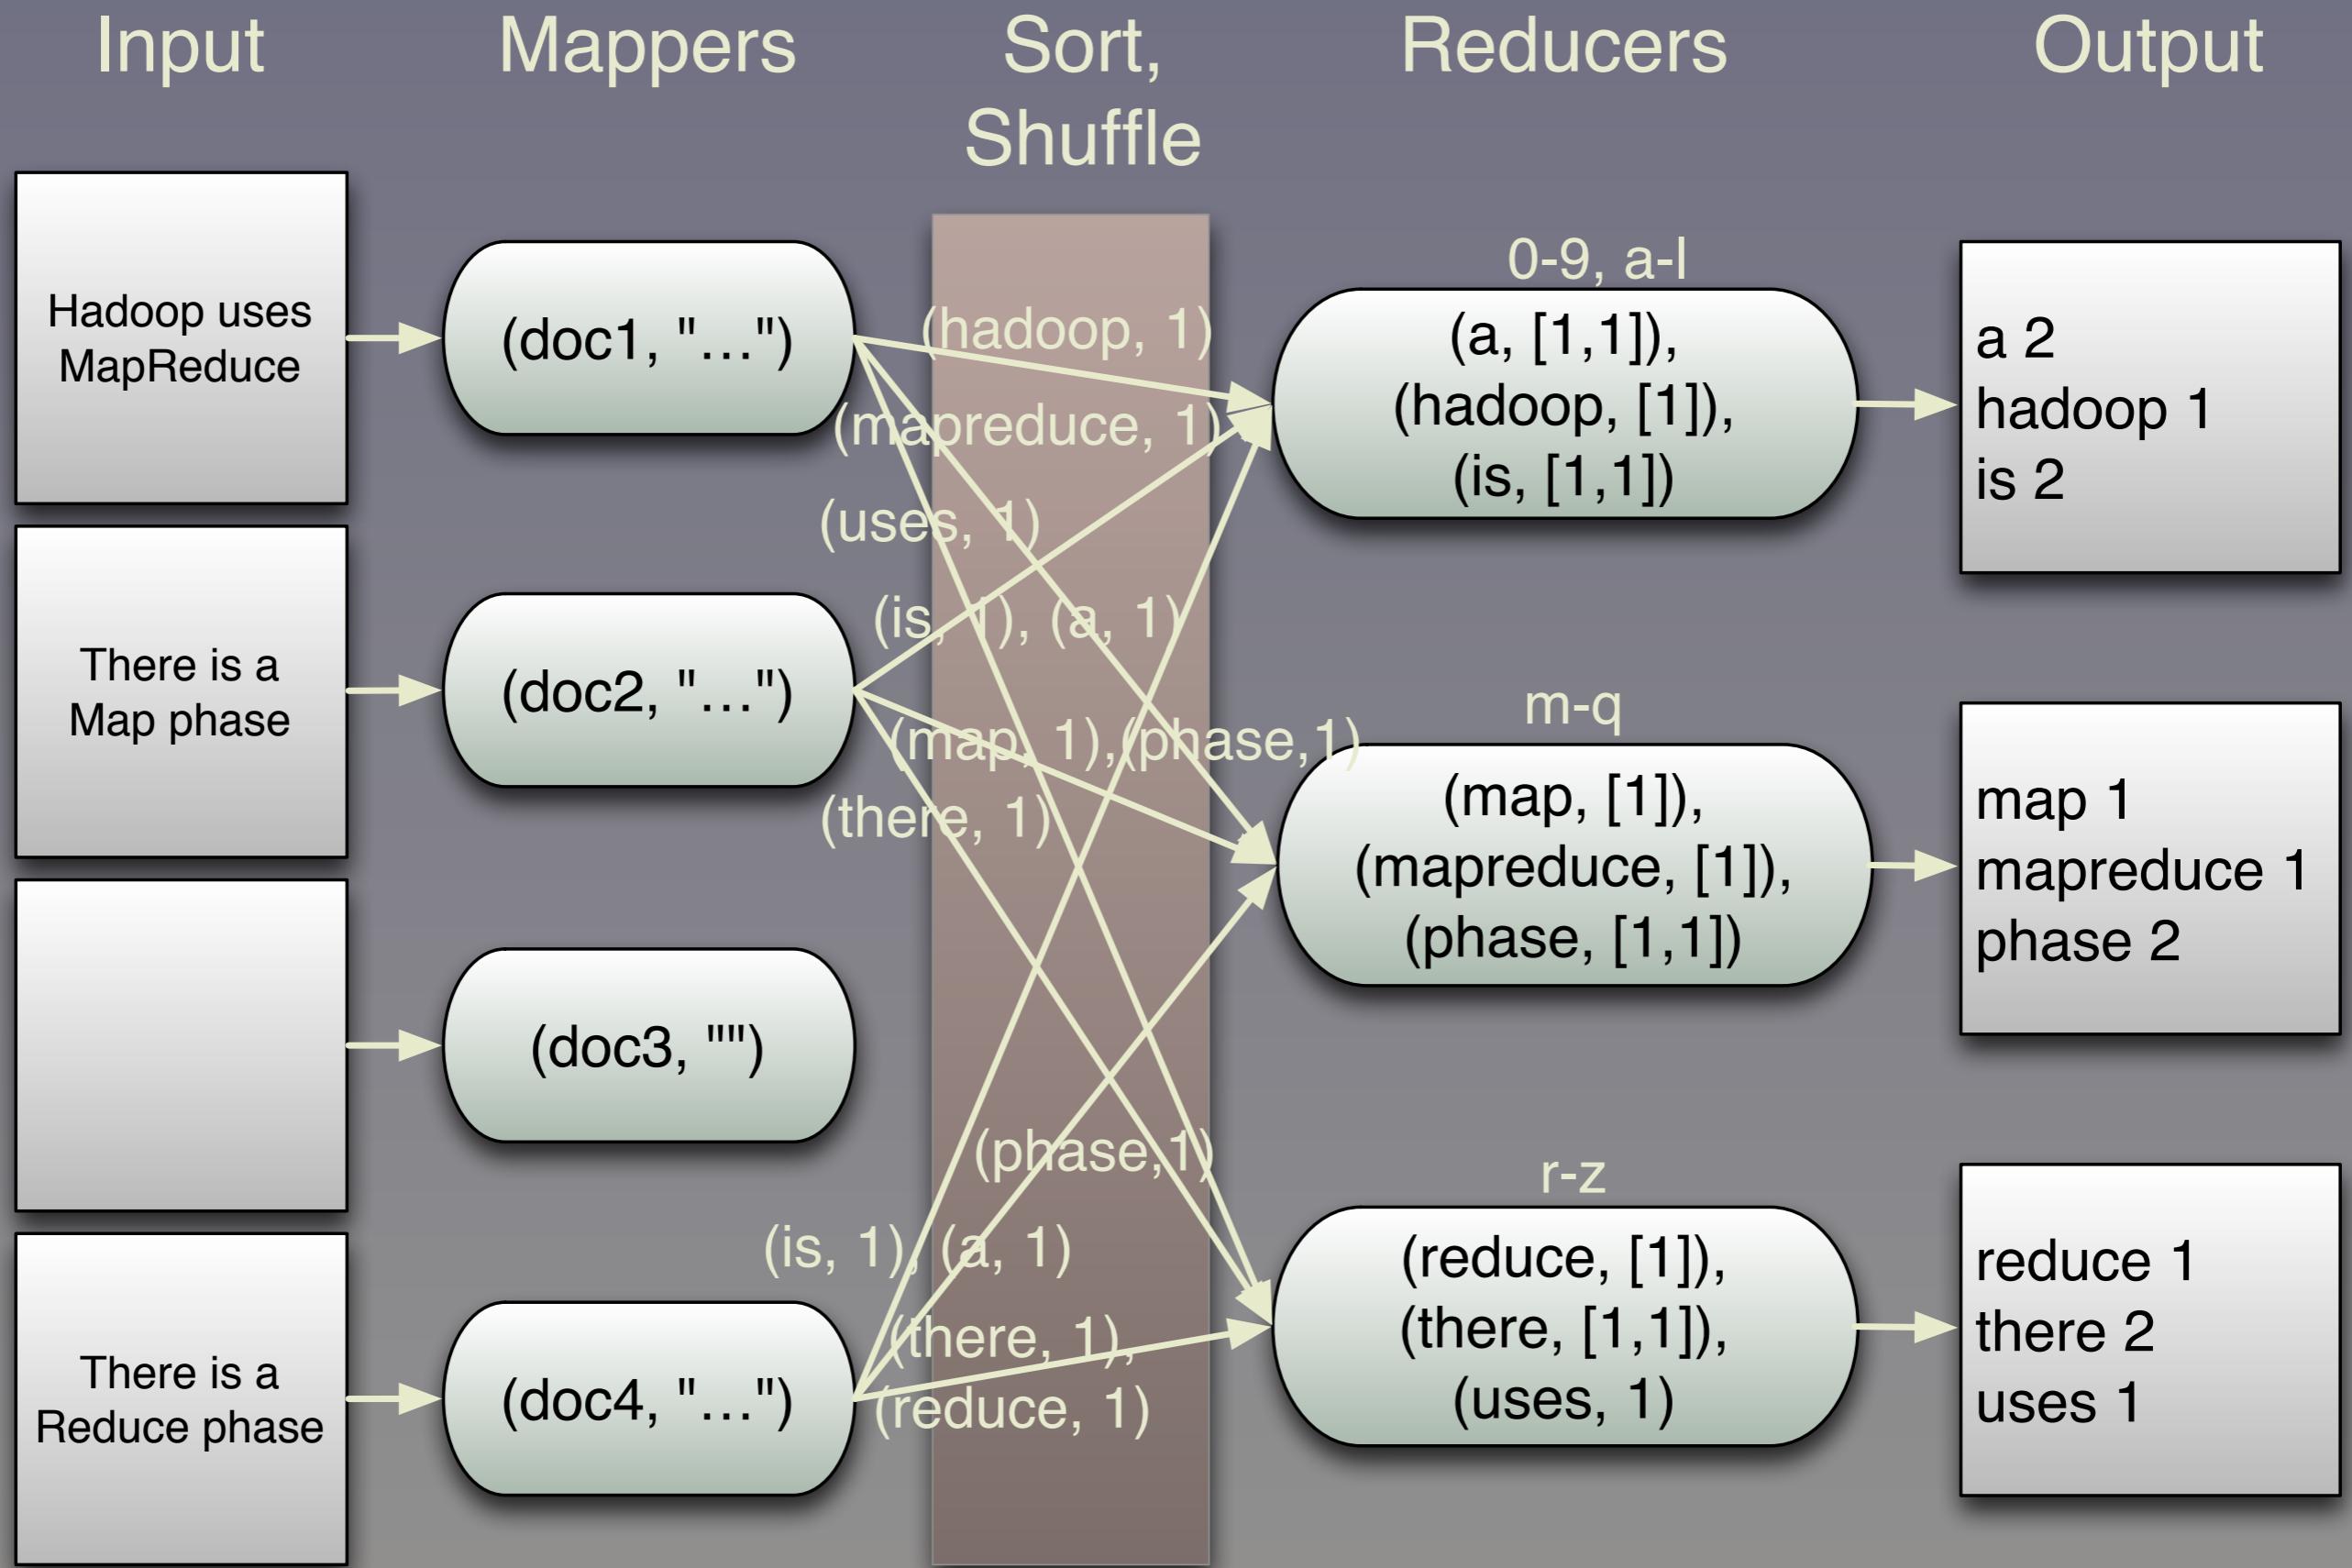
Note how we partitioned the reducers, by first letter of the keys. (By default, MR just hashes the keys and distributes them modulo # of reducers.)

Input Mappers Sort, Shuffle Reducers



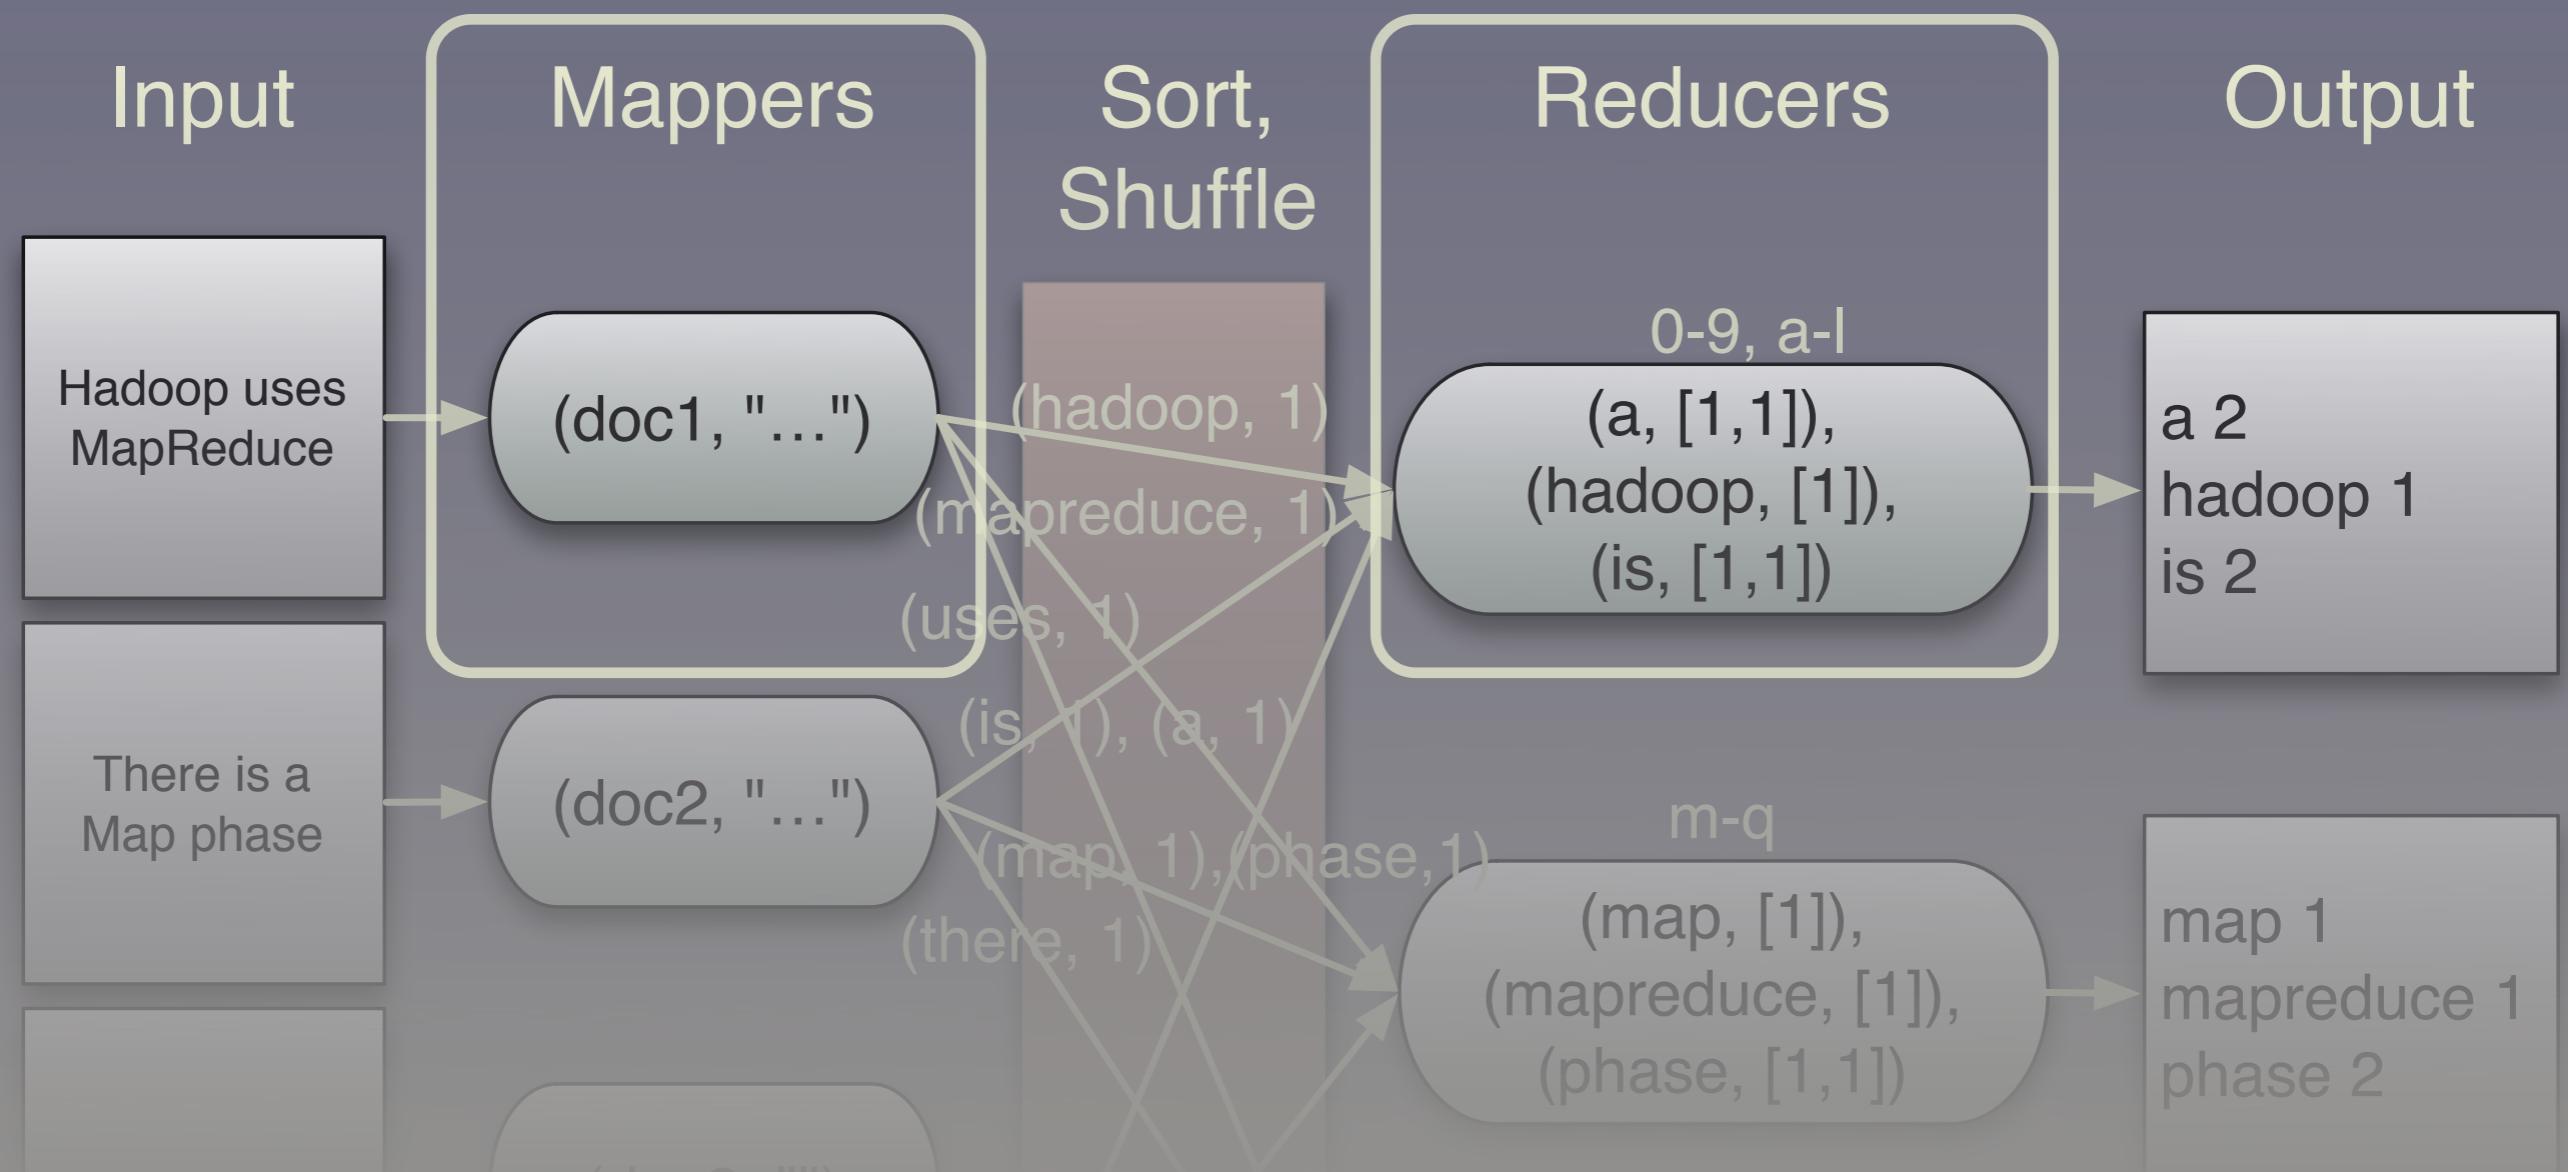
Tuesday, June 19, 12

Our reducers are passed each key and a collection of all the values for that key. (The MR framework creates this collection for us.)



Tuesday, June 19, 12

The final view of the WordCount process flow. The reducer just sums the counts and writes the output. The output files contain one line for each key (the word) and value (the count), assuming we're using text output. The choice of delimiter between key and value is up to you, but tab is common. (We'll discuss options as we go.)



Map:

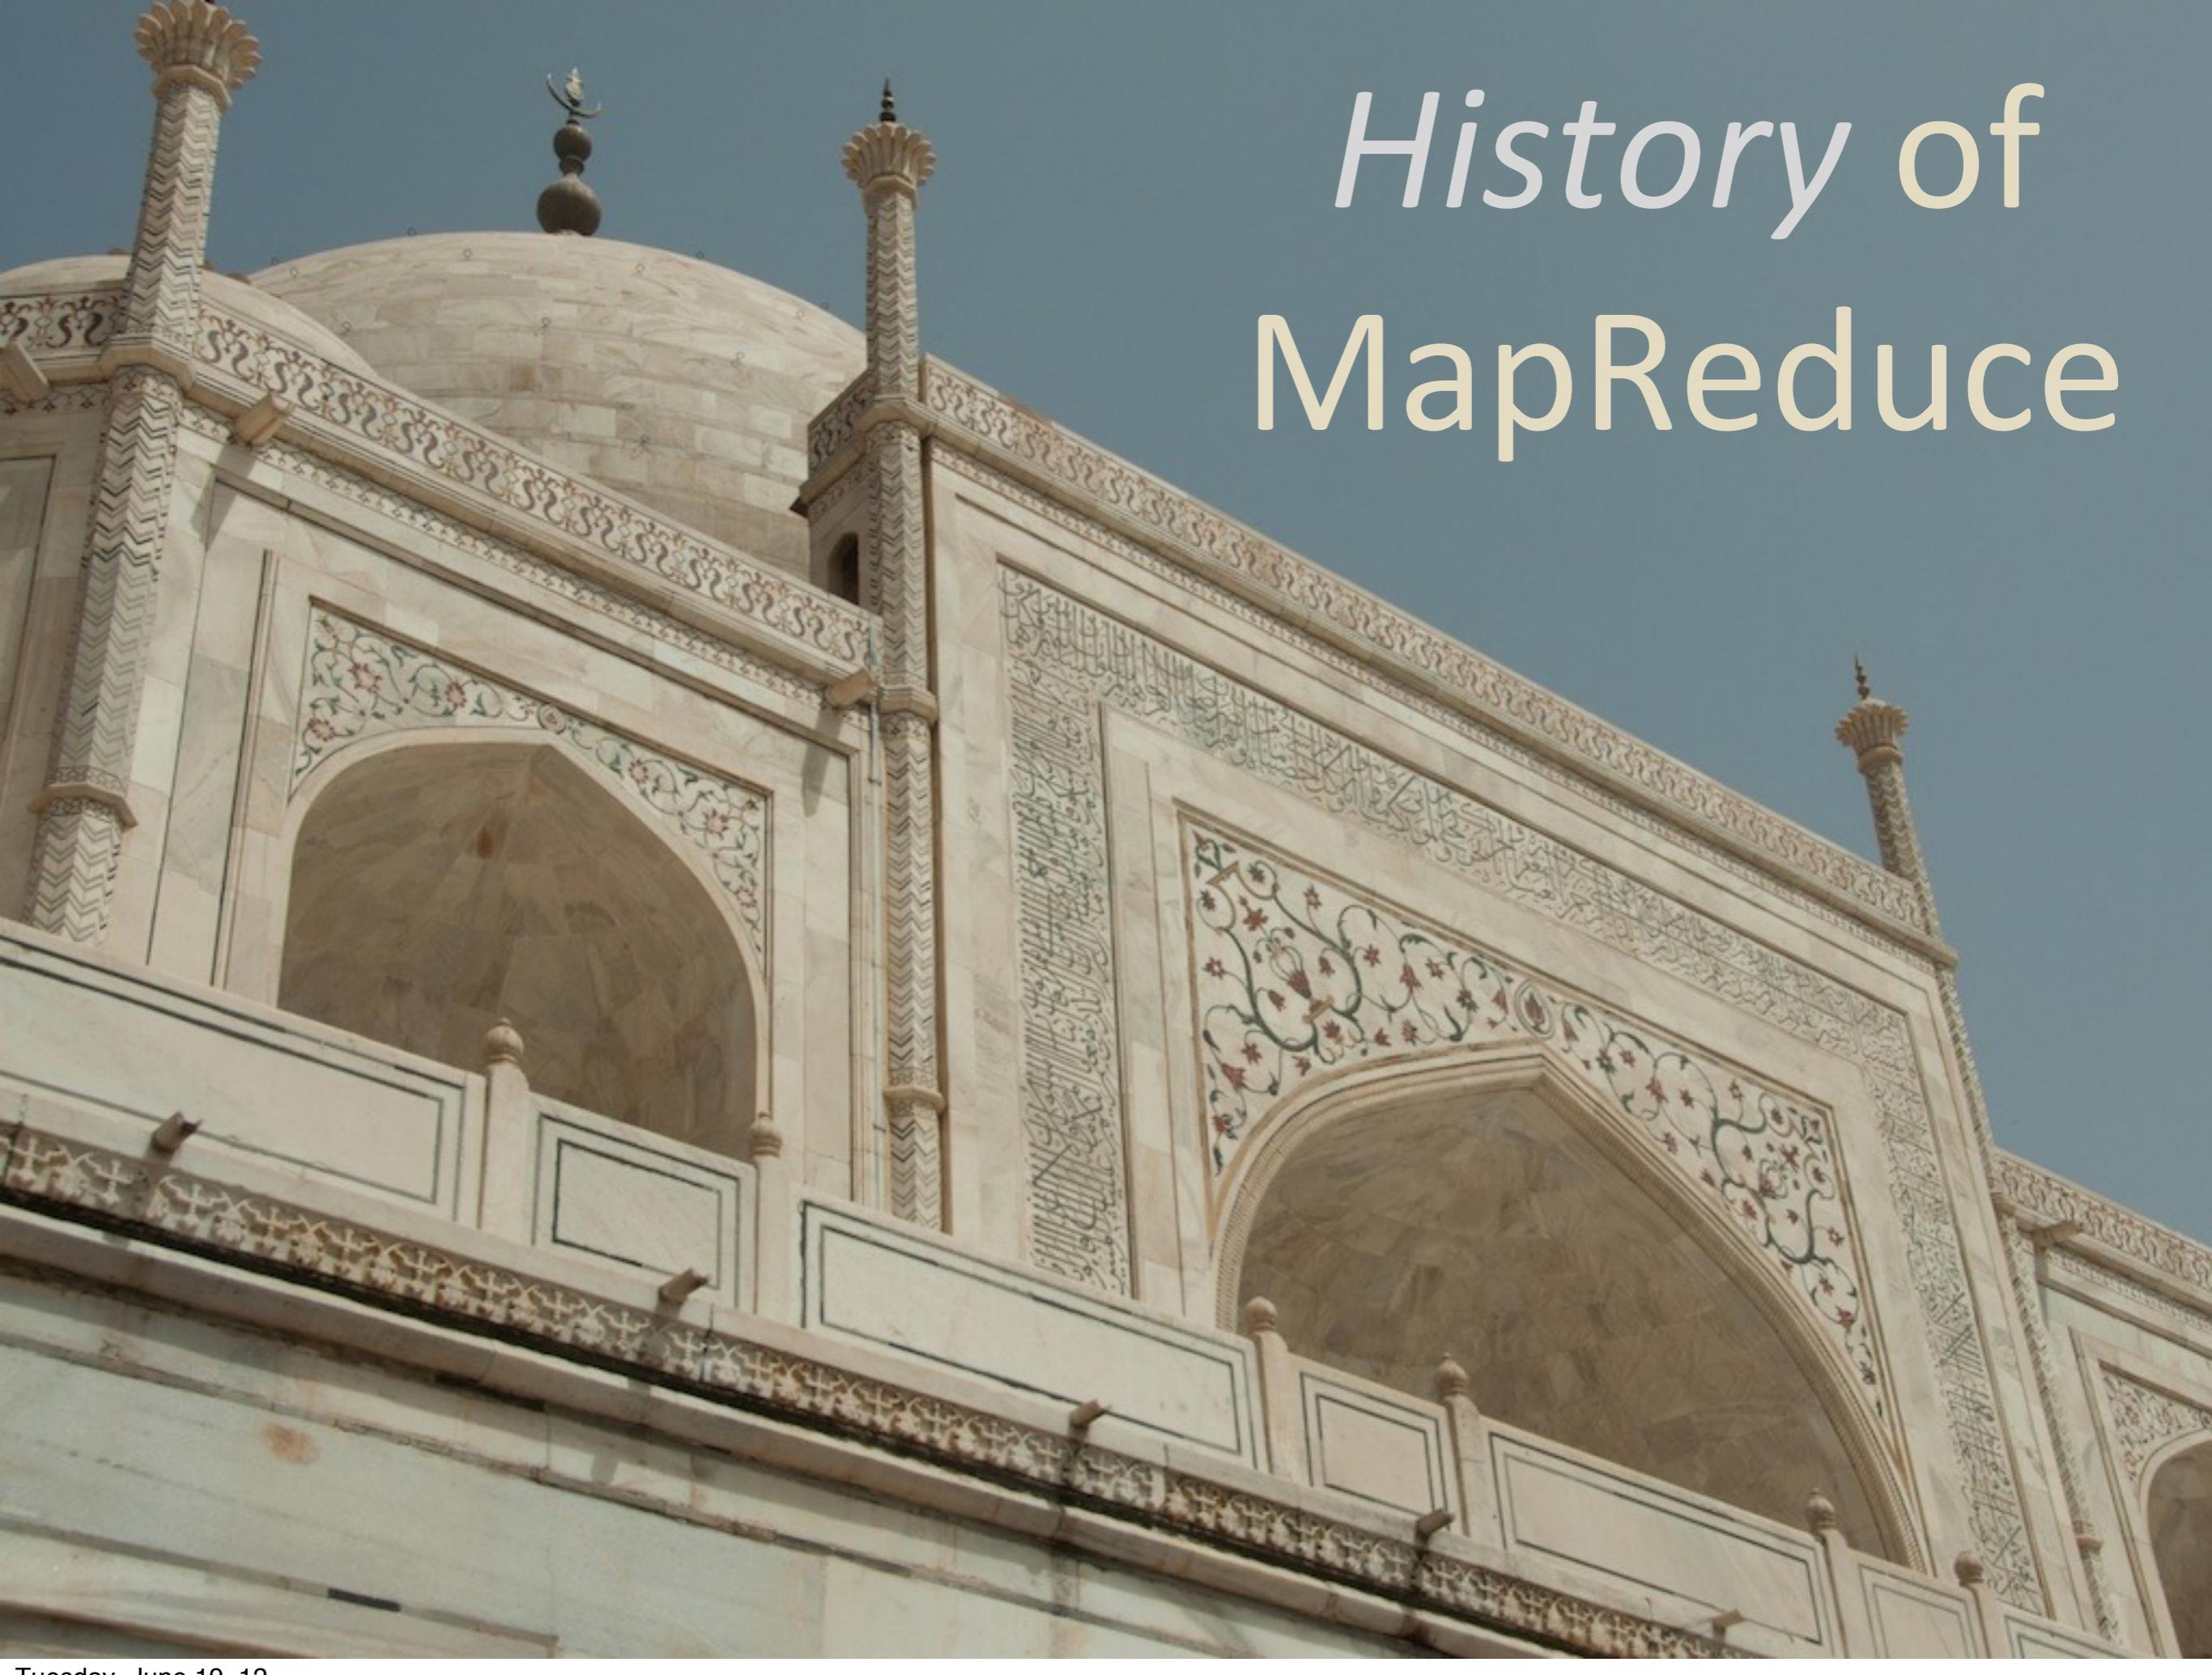
- Transform *one* input to *0-N* outputs.

Reduce:

- Collect *multiple* inputs into one output.

Tuesday, June 19, 12

To recap, a “map” transforms one input to one output, but this is generalized in MapReduce to be one to 0-N. The output key-value pairs are distributed to reducers. The “reduce” collects together multiple inputs with the same key into

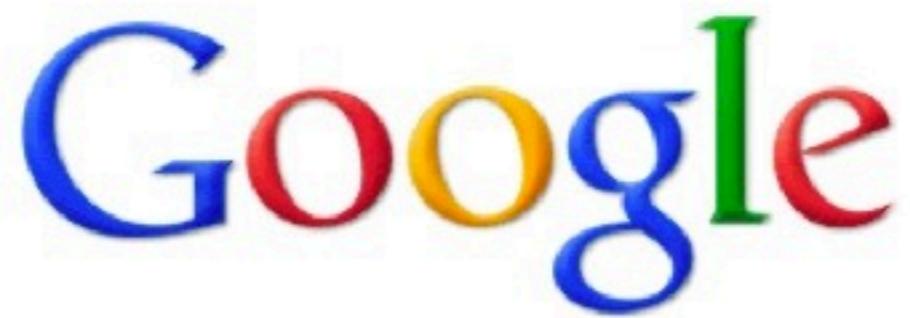


History of MapReduce

Tuesday, June 19, 12

Let's review where MapReduce came from and its best-known, open-source incarnation, Hadoop.

How would you *index the web?*

A screenshot of a Google search interface. The search bar contains the partial query "What is the meanin". Below the bar, a list of suggested completions appears, all starting with "what is the meaning of":

- what is the meaning of life
- what is the meaning of life 42
- what is the meaning of pumped up kicks
- what is the meaning of halloween
- what is the meaning of love
- what is the meaning of labor day
- what is the meaning of slope
- what is the meaning of homecoming
- what is the meaning of my last name
- what is the meaning of a promise ring

At the bottom of the interface are two buttons: "Google Search" and "I'm Feeling Lucky".

How would you *index the web?*

Google

what is the meaning of life

Search

About 48,900,000 results (0.26 seconds)

Everything

Images

Maps

Videos

News

Shopping

[Meaning of life - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Meaning_of_life +1
The **meaning of life** constitutes a philosophical question concerning the purpose and significance of life or existence in general. This concept can be expressed ...
[Questions - Western philosophical perspectives - East Asian philosophy](#)

[The Meaning of Life](#)
users.aristotle.net/~diogenes/meaning1.htm +1
Why do you want to know the **meaning of life**? Often people ask this question when they really want the answer to some other question. Let's try and get those ...

»

You ask a *phrase* and
the *search engine* finds
the *best match* in
billions of web pages.

Actually, Google
computes the *index*
that *maps* terms to
pages in *advance*.

Google's famous *Page Rank* algorithm.

Google File System

for Storage

2003

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung

Google*

ABSTRACT

We have designed and implemented the Google File System, a scalable distributed file system for large distributed data-intensive applications. It provides fault tolerance while running on inexpensive commodity hardware, and it delivers high aggregate performance to a large number of clients.

While sharing many of the same goals as previous distributed file systems, our design has been driven by observations of our application workloads and technological envi-

1. INTRODUCTION

We have designed and implemented the Google File System (GFS) to meet the rapidly growing demands of Google's data processing needs. GFS shares many of the same goals as previous distributed file systems such as performance, scalability, reliability, and availability. However, its design has been driven by key observations of our application workloads and technological environment, both current and anticipated, that reflect a marked departure from some earlier

Tuesday, June 19, 12

A distributed file system provides horizontal scalability and resiliency when file blocks are duplicated around the cluster.

MapReduce for Computation

2004

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to par-

Tuesday, June 19, 12

A distributed file system provides horizontal scalability and resiliency when file blocks are duplicated around the cluster.

About this time, *Doug Cutting*, the creator of *Lucene* was working on *Nutch*...

He implemented clean-
room versions of
MapReduce and *GFS*...

By 2006 , they became
part of a separate
Apache project,
called *Hadoop*.



Tuesday, June 19, 12

The name comes from a toy, stuffed elephant that Cutting's son owned at the time.

Benefits of *MapReduce*



Tuesday, June 19, 12

The best way to
approach *Big Data* is to
scale *Horizontally*.

MapReduce Design Goals

Maximize Disk I/O!!

Oh, and run on *server-class, commodity hardware.*

Great for *Batch Mode*
Data Crunching.

MapReduce and Its *Discontents*



Tuesday, June 19, 12

Is MapReduce the end of the story? Does it meet all our needs?

#1

The *Hadoop Java API*
is too *verbose*
& *object oriented.*

Tuesday, June 19, 12

Even word count is not “obvious”. When you get to fancier stuff like joins, group-bys, etc., the mapping is not trivial at all.

#2

It's hard *implementing*
many *Algorithms*
in *MapReduce*.

Tuesday, June 19, 12

Even word count is not “obvious”. When you get to fancier stuff like joins, group-bys, etc., the mapping is not trivial at all.

#3

Some algorithms have
suboptimal performance
in MapReduce.

For example, *Graph algorithms*.

Tuesday, June 19, 12

While general purpose, some algorithms are not as fast as they could be, such as graph traversal algorithms (like Page Rank!!), because typically each traversal of an “arc” is a computation step. Hadoop is bad because every such ‘step’ is a separate process with a new JVM instance (by default). Also there is a fair amount of overhead moving data between mappers and reducers. (Google’s implementation may be more efficient.)

#4

What about
event processing
("real-time")?

Tuesday, June 19, 12

Even Google now does not want to wait to compute a new index of the web. As soon as its web crawlers detect changes, Google would like to update its index as soon as possible.

The background of the slide features a wide-angle photograph of a desert landscape at sunset or sunrise. In the foreground, there's a dark, flat area with sparse vegetation. Behind it, several large, layered red rock mesas rise against a sky filled with scattered, warm-colored clouds.

The *Hadoop* Java API

#1 & #2

Tuesday, June 19, 12

Let's explore these issues, starting with the Hadoop Java API, but I'll show you Scala code; see my GitHub `scala-hadoop` project for the full code examples.

```
import org.apache.hadoop.io._  
import org.apache.hadoop.mapred._  
import java.util.StringTokenizer  
  
object SimpleWordCount {  
  
    val one = new IntWritable(1)  
    val word = new Text      // Value will be set in a non-thread-safe way!  
  
    class WCMapper extends MapReduceBase with Mapper[LongWritable, Text, Text, IntWritable] {  
  
        def map(key: LongWritable, valueDocContents: Text,  
               output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {  
            val tokens = valueDocContents.toString.split("\\s+")  
            for (wordString <- tokens) {  
                if (wordString.length > 0) {  
                    word.set(wordString.toLowerCase)  
                    output.collect(word, one)  
                }  
            }  
        }  
    }  
  
    class Reduce extends MapReduceBase with Reducer[Text, IntWritable, Text, IntWritable] {  
  
        def reduce(keyWord: Text, valuesCounts: java.util.Iterator[IntWritable],  
                  output: OutputCollector[Text, IntWritable], reporter: Reporter): Unit = {  
            var totalCount = 0  
            while (valuesCounts.hasNext) {  
                totalCount += valuesCounts.next.get  
            }  
            output.collect(keyWord, new IntWritable(totalCount))  
        }  
    }  
}
```

Tuesday, June 19, 12

This is intentionally too small to read. The algorithm is simple, but the framework is in your face. Note all the green types floating around and relatively few yellow methods implementing actual operations. Still, I've omitted many boilerplate details for configuring and running the job. This is just the "core" MapReduce code. In fact, Word Count is not too bad, but when you get to more complex algorithms, even conceptually simple things like relational-style joins, code in this API gets complex and tedious very fast.
I'm using Scala to call the Java API. If this were Java code, it would be more verbose.

The API is *invasive*
and it *obscures*
the application *logic*...

A MapReduce Assembly Language.

Tuesday, June 19, 12

There is a very poor “separation of concerns”

How can we do better?

Decouple *logic*
from *infrastructure*.

Tuesday, June 19, 12

There is a very poor “separation of concerns”

How can we do better?

*Use Functional
Programming!!*

Tuesday, June 19, 12

This is data we're talking about, so we should be using mathematics. FP is as close as programming comes to mathematics!

Using Crunch (Java)

(Addressing #1 and #2)



Tuesday, June 19, 12

Let's start by looking at options that improve the "MapReduce experience", but still work within that framework, but with a better API. Also, these options address #2, the difficulty of translating algorithms to MapReduce.

Crunch is a Java library that provides a higher-level abstraction for data computations and flows on top of MapReduce, inspired by a paper from Google researchers on an internal project called FlumeJava.

See <https://github.com/cloudera/crunch>.

Crunch Concepts

*Distributed Collections,
Parallel Operations,
and Pipelines.*

Tuesday, June 19, 12

The collections include unordered sets and maps. Operations on these collections can be performed in parallel, including filtering, grouping, etc. Pipelines abstract over data processing flows. Note it was necessary to implement functional collection abstractions missing from Java.

Java

```
import com.cloudera.crunch.*;
import org.apache.hadoop.*;
...
public class WordCount extends Configured implements Tool, Serializable {
    public int run(String[] args) throws Exception {
        Pipeline pipeline = new MRPipeline(WordCount.class, getConf());
        PCollection<String> lines = pipeline.readTextFile(args[0]);
        PCollection<String> words = lines.parallelDo(new DoFn<String, String>() {
            public void process(String line, Emitter<String> emitter) {
                for (String word : line.split("\\s+")) {
                    emitter.emit(word);
                }
            }
        }, Writables.strings());
        PTable<String, Long> counts = Aggregate.count(words);
        pipeline.writeTextFile(counts, args[1]);
        pipeline.done();
        return 0;
    }
}
```

... still busy

Tuesday, June 19, 12

This is back to Java (rather than Scala calling a Java API). I omitted the setup and “main” again, but that code is a lot smaller than for generic MapReduce. It’s a definite improvement, especially for more sophisticated algorithms. Crunch (and FlumeJava), as well as similar toolkits like Cascading, are steps in the right direction, but there is still 1) lots of boilerplate, 2) functional ideas crying to escape the Java world view, and poor support in the MapReduce API (like no Tuple type) that harms Crunch, etc. Still, note that relatively heavy amount of type information compared to methods (operations). It would improve some if I wrote this code in Scala, thereby exploiting type inference, but not by a lot.

Using Scrunch (Scala)



Tuesday, June 19, 12

Scrunch is a scala DSL around Crunch writing by the same developers at Cloudera and also included in the Crunch distro. It uses type classes and similar constructs to provide wrap the Crunch classes and MR classes with REAL map, flatMap, etc. functionality.

```
import com.cloudera.crunch._  
import com.cloudera.scrunch._
```

```
...
```

Scala

```
class ScrunchWordCount {  
    def wordCount(inputFile: String,  
                  outputFile: String) = {  
  
        val pipeline = new Pipeline[WordCountExample]  
        pipeline.read(from.textFile(inputFile))  
            .flatMap(_.toLowerCase.split("\\\\W+"))  
            .filter(!_isEmpty())  
            .count  
            .write(to.textFile(outputFile)) // Word counts  
            .map((w, c) => (w.slice(0, 1), c))  
            .groupByKey.combine(v => v.sum)  
            .materialize  
        pipeline.done  
    }  
}
```

```
object ScrunchWordCount {  
    def main(args: Array[String]) = {  
        new ScrunchWordCount.wordCount(args(0), args(1))  
    }  
}
```

Is nice!

Tuesday, June 19, 12

(Back to Scala) I cheated; I'm showing you the WHOLE program, "main" and all. Not only is the size significantly smaller and more concise still, but note the "builder" style notation, which intuitive lays out the data flow required. There is must less green - fewer types are explicitly tossed about, and not just because Scala does implicit typing. In contrast, there is more yellow - function calls showing the sequence of operations that more naturally represent the real "business logic", i.e., the data flow that is Word Count!

Also, fewer comments are required to help you sort out what's going on. Once you understand the meaning of the individual, high-reusable operations like flatMap and filter, you can construct arbitrarily-complex transformations and calculations with relative ease.

A *DSL* of *mathematical* operations.

Tuesday, June 19, 12

Domain-Specific Language that hides the “assembly language” and exposes the most appropriate operations for working with data, various mathematical abstractions...

Data \Leftrightarrow Mathematics

*∴ use functional
programming!*

Tuesday, June 19, 12

... and functional programming models mathematical operations ...

Data \leftrightarrow Mathematics

e.g., *map, flat map, fold, reduce, group by...*

Tuesday, June 19, 12

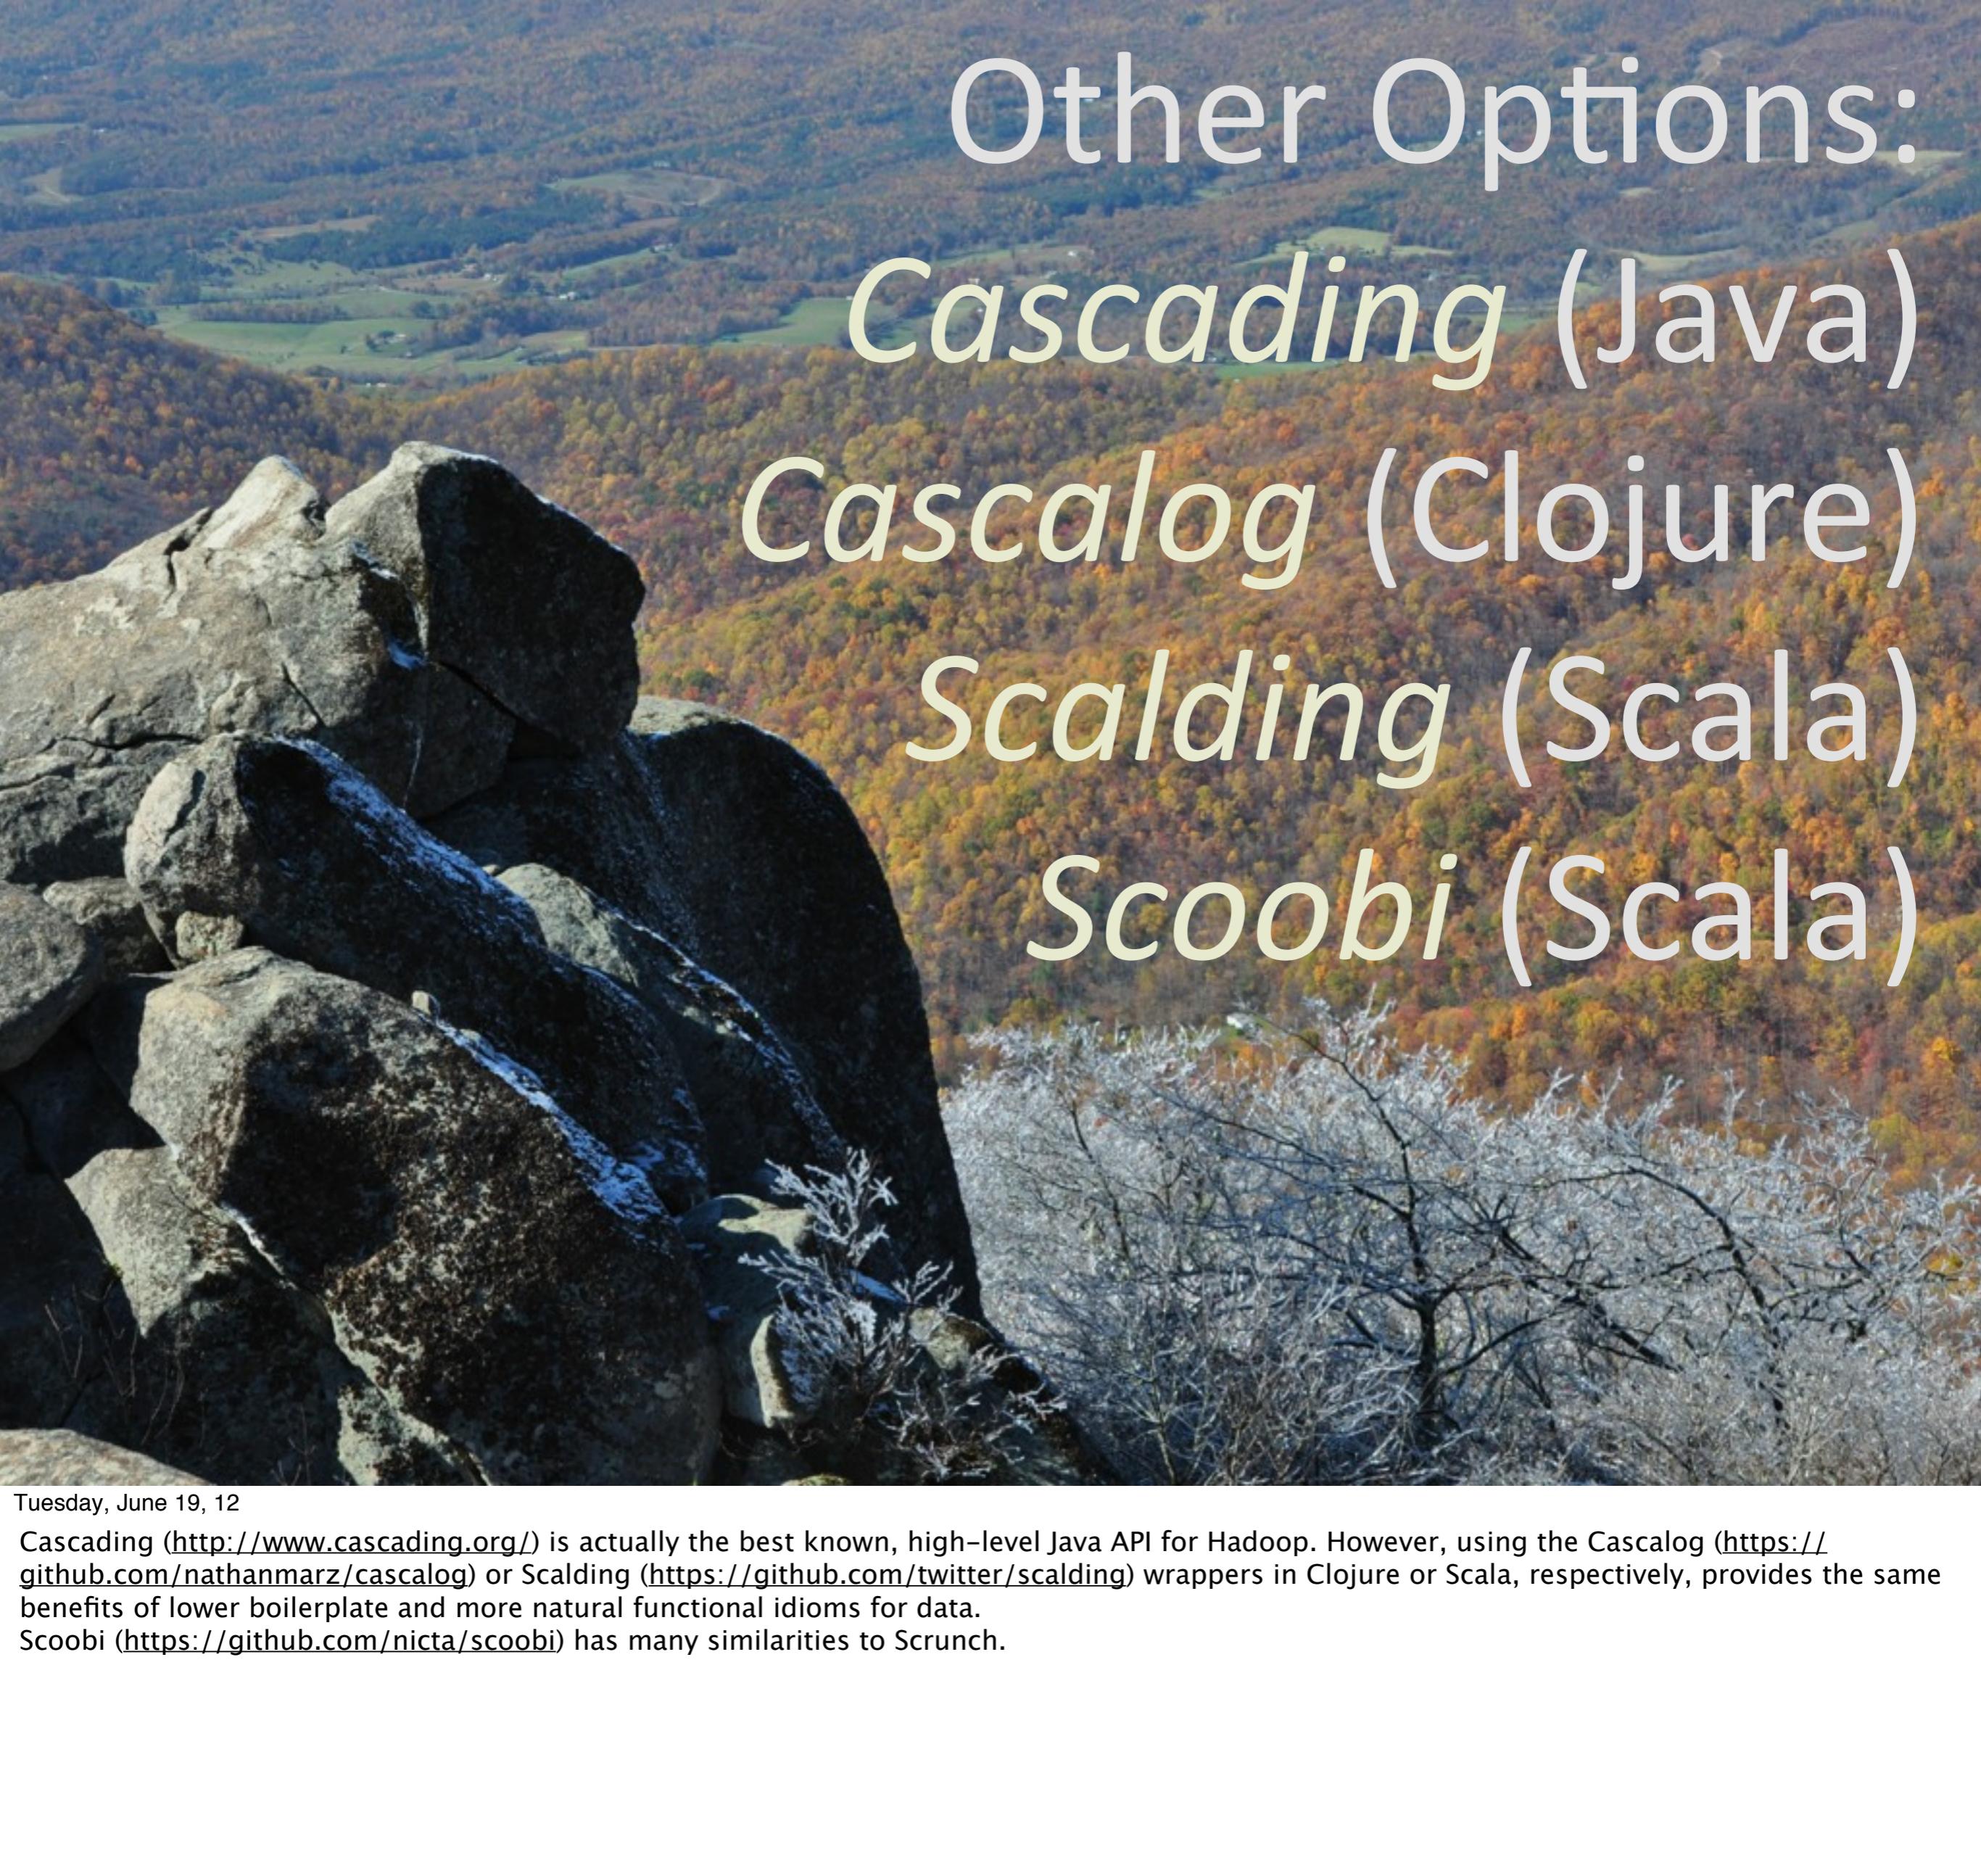
... such as map, flatmap, fold, reduce, ...

Keep in mind that “MapReduce” originated from functional concepts, just bloated up...

There are other
imperative-OO flaws
in Hadoop...

Tuesday, June 19, 12

Like reuse of objects (premature optimization), few natural key-value types (e.g., no tuple type available). I see a lot of Java bloat and poor modularity that make extensions and optimizations hard to implement.



Other Options:

- Cascading* (Java)
- Cascalog* (Clojure)
- Scalding* (Scala)
- Scoobi* (Scala)

Tuesday, June 19, 12

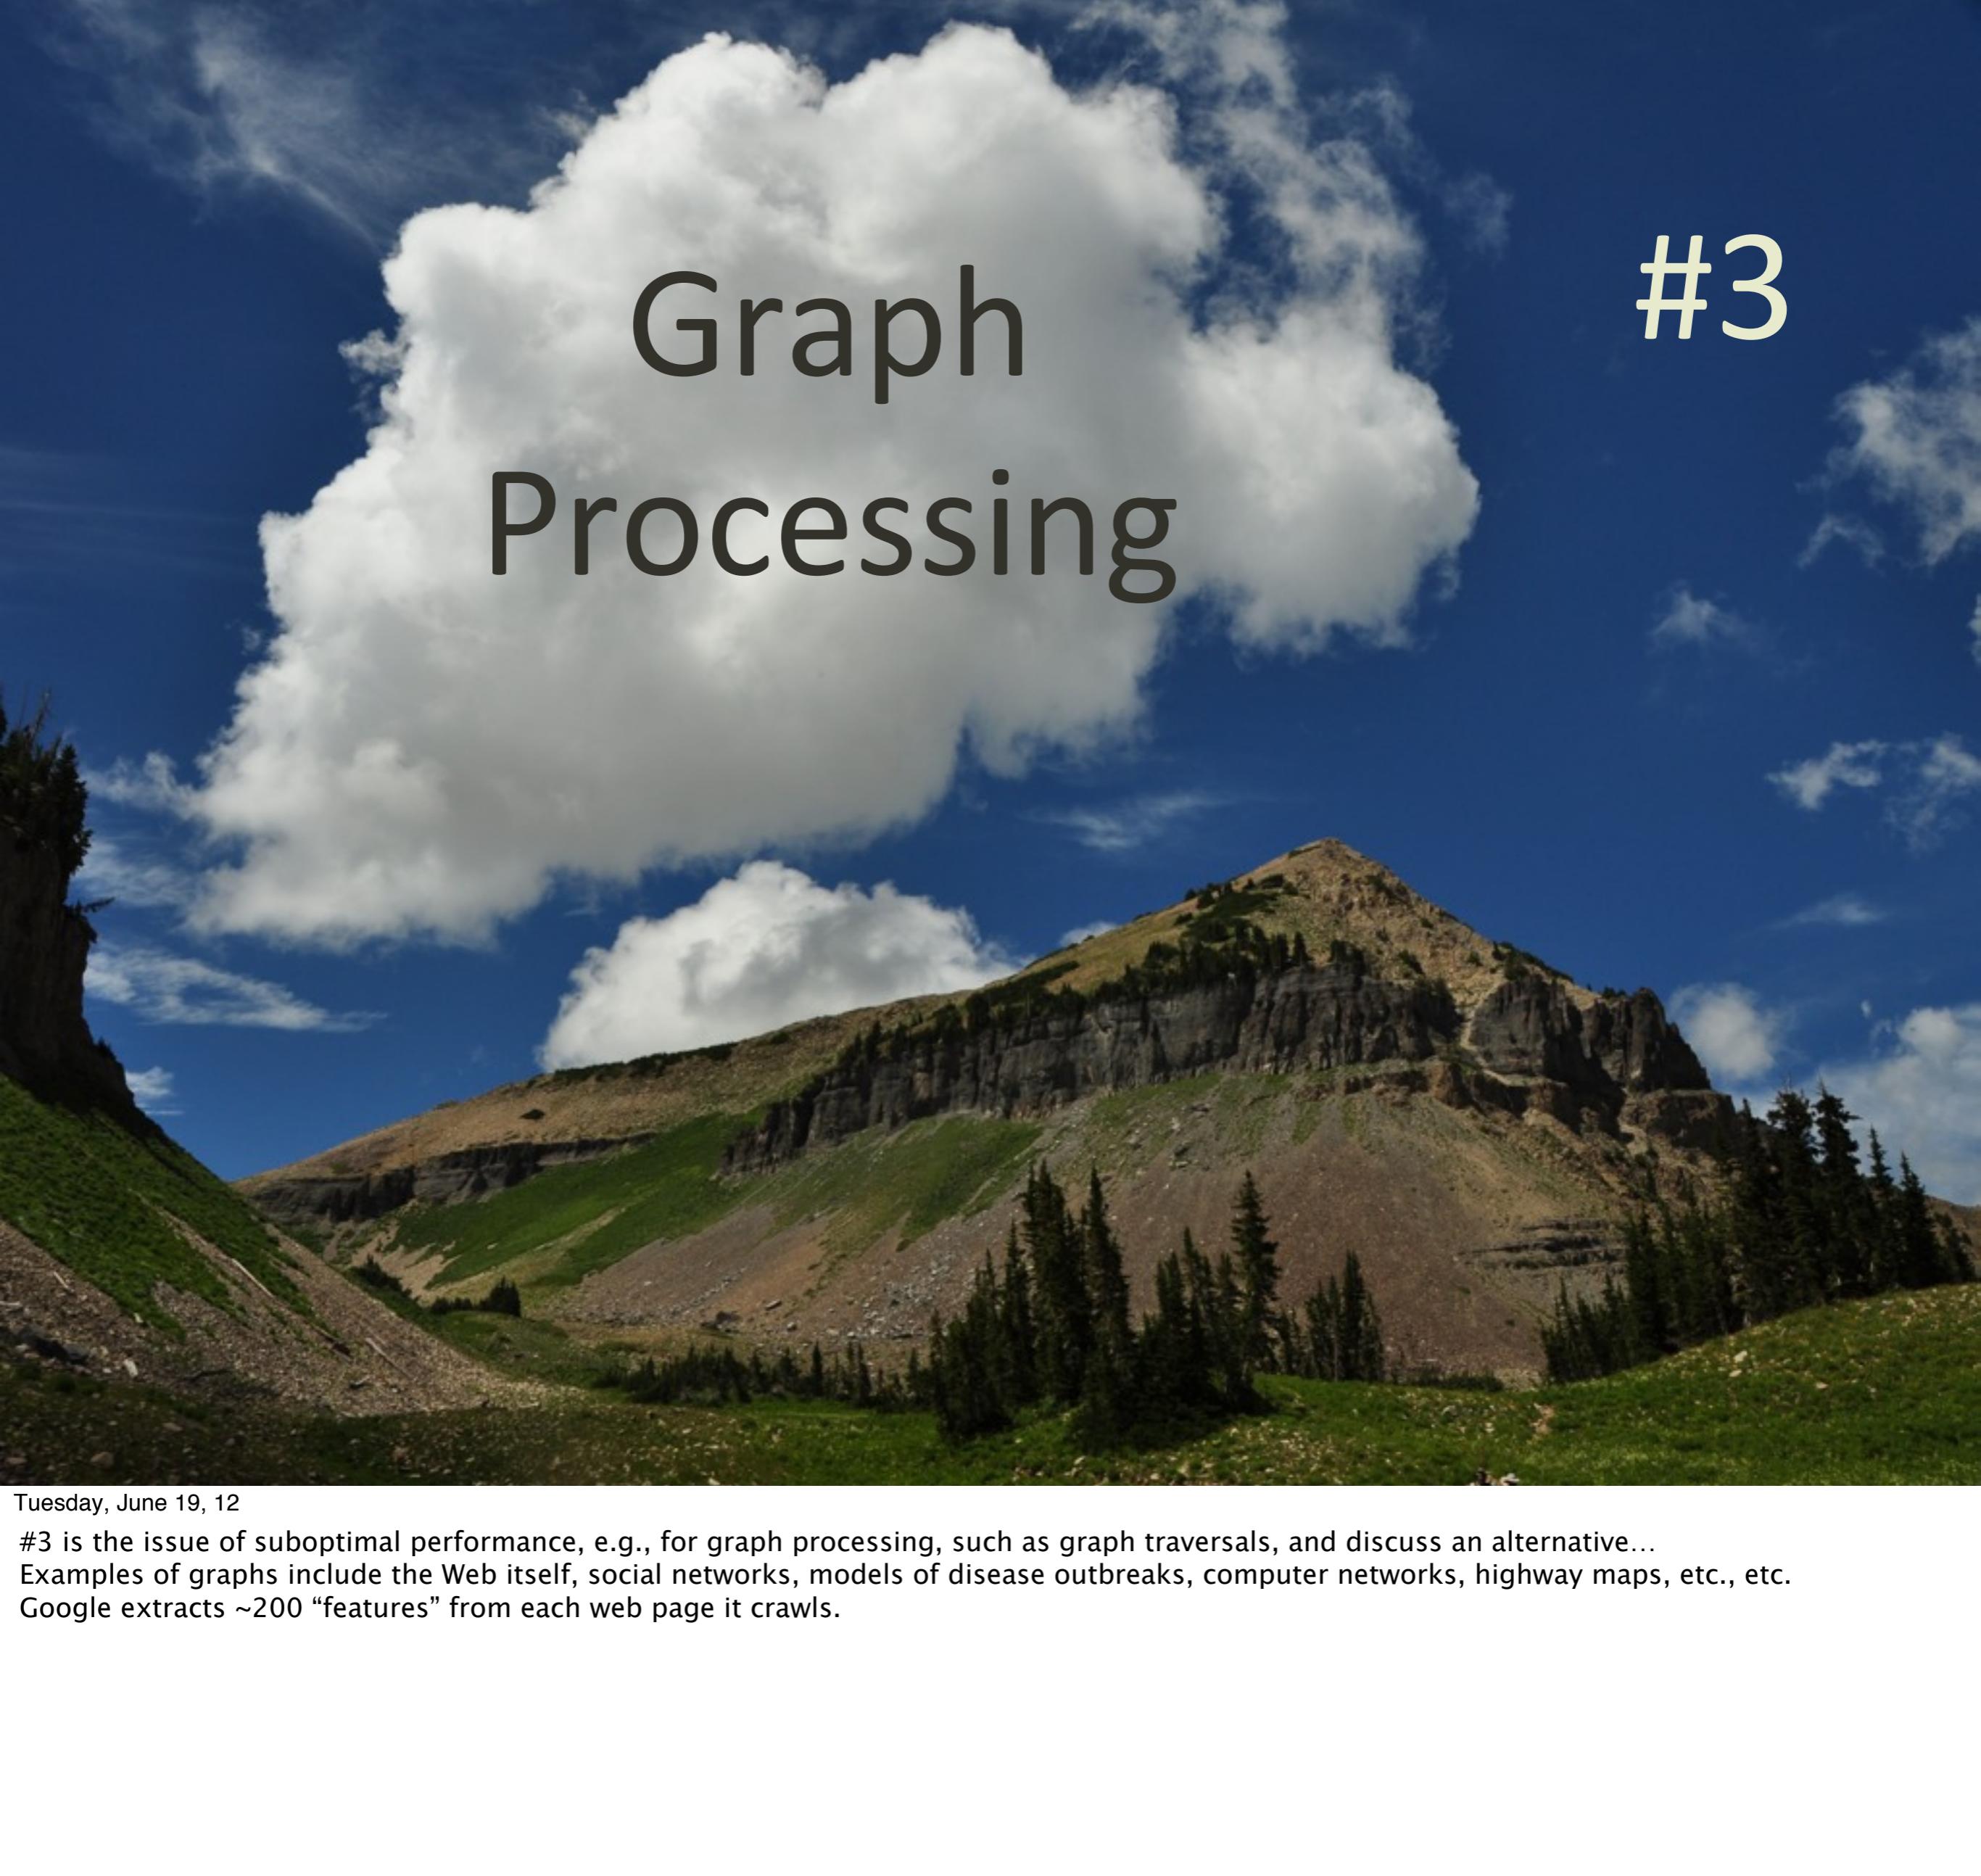
Cascading (<http://www.cascading.org/>) is actually the best known, high-level Java API for Hadoop. However, using the Cascalog (<https://github.com/nathanmarz/cascalog>) or Scalding (<https://github.com/twitter/scalding>) wrappers in Clojure or Scala, respectively, provides the same benefits of lower boilerplate and more natural functional idioms for data.
Scoobi (<https://github.com/nicta/scoobi>) has many similarities to Scrunch.

The background of the slide is a scenic landscape of a mountain range during autumn. The trees on the hillsides are a mix of green, yellow, and orange. In the foreground, there are large, dark, craggy rocks. A few bare branches of a tree are visible in the lower right corner.

Other Options: *Hive and Pig*

Tuesday, June 19, 12

Hive and Pig are popular parts of the Hadoop ecosystem, but they are limited. Hive is a SQL query language, which is actually pretty impressive, but not full-featured, Turing-complete language. Neither is Pig, which is a data flow language. (To be fair, both can be extended with code written in Java or other languages, but you avoid that with the Java/Clojure/Scala toolkits.)



Graph Processing

#3

Tuesday, June 19, 12

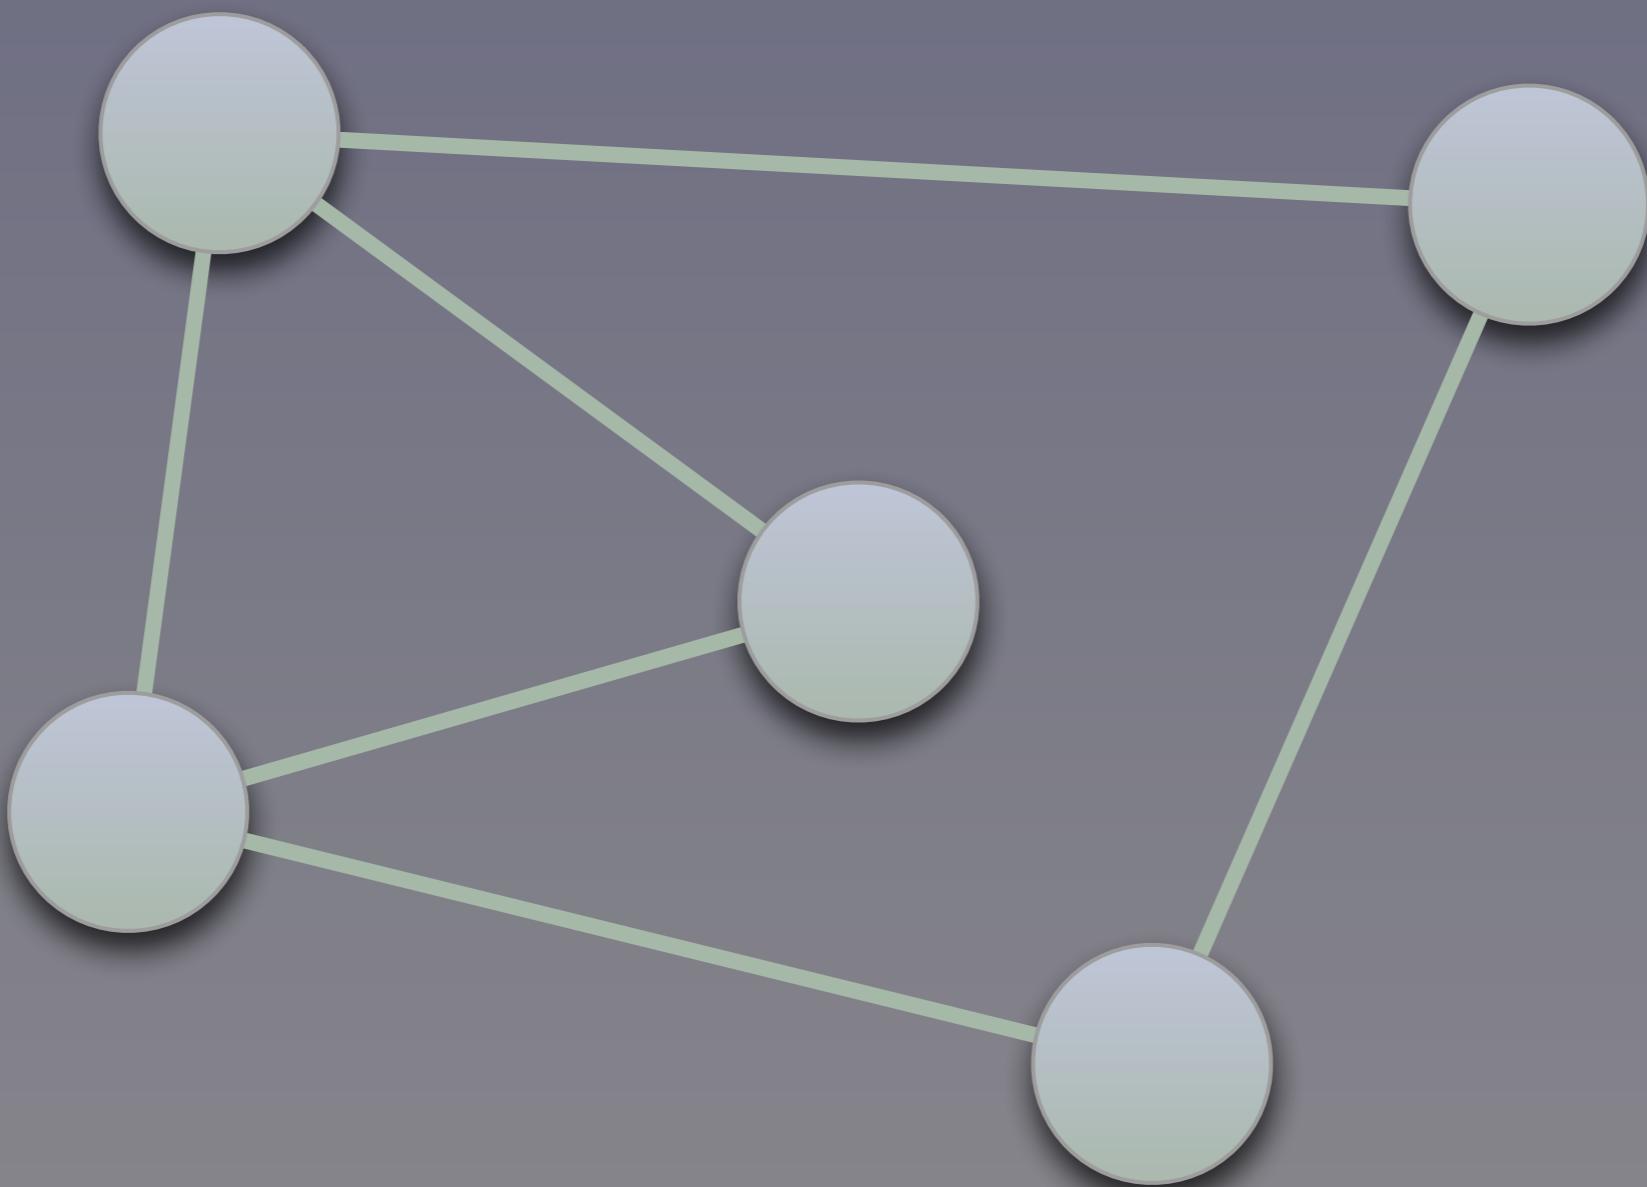
#3 is the issue of suboptimal performance, e.g., for graph processing, such as graph traversals, and discuss an alternative...
Examples of graphs include the Web itself, social networks, models of disease outbreaks, computer networks, highway maps, etc., etc.
Google extracts ~200 “features” from each web page it crawls.

*Google is actually
moving away from
Page Rank on
MapReduce to Pregel.*

Tuesday, June 19, 12

see <http://googleresearch.blogspot.com/2009/06/large-scale-graph-computing-at-google.html>

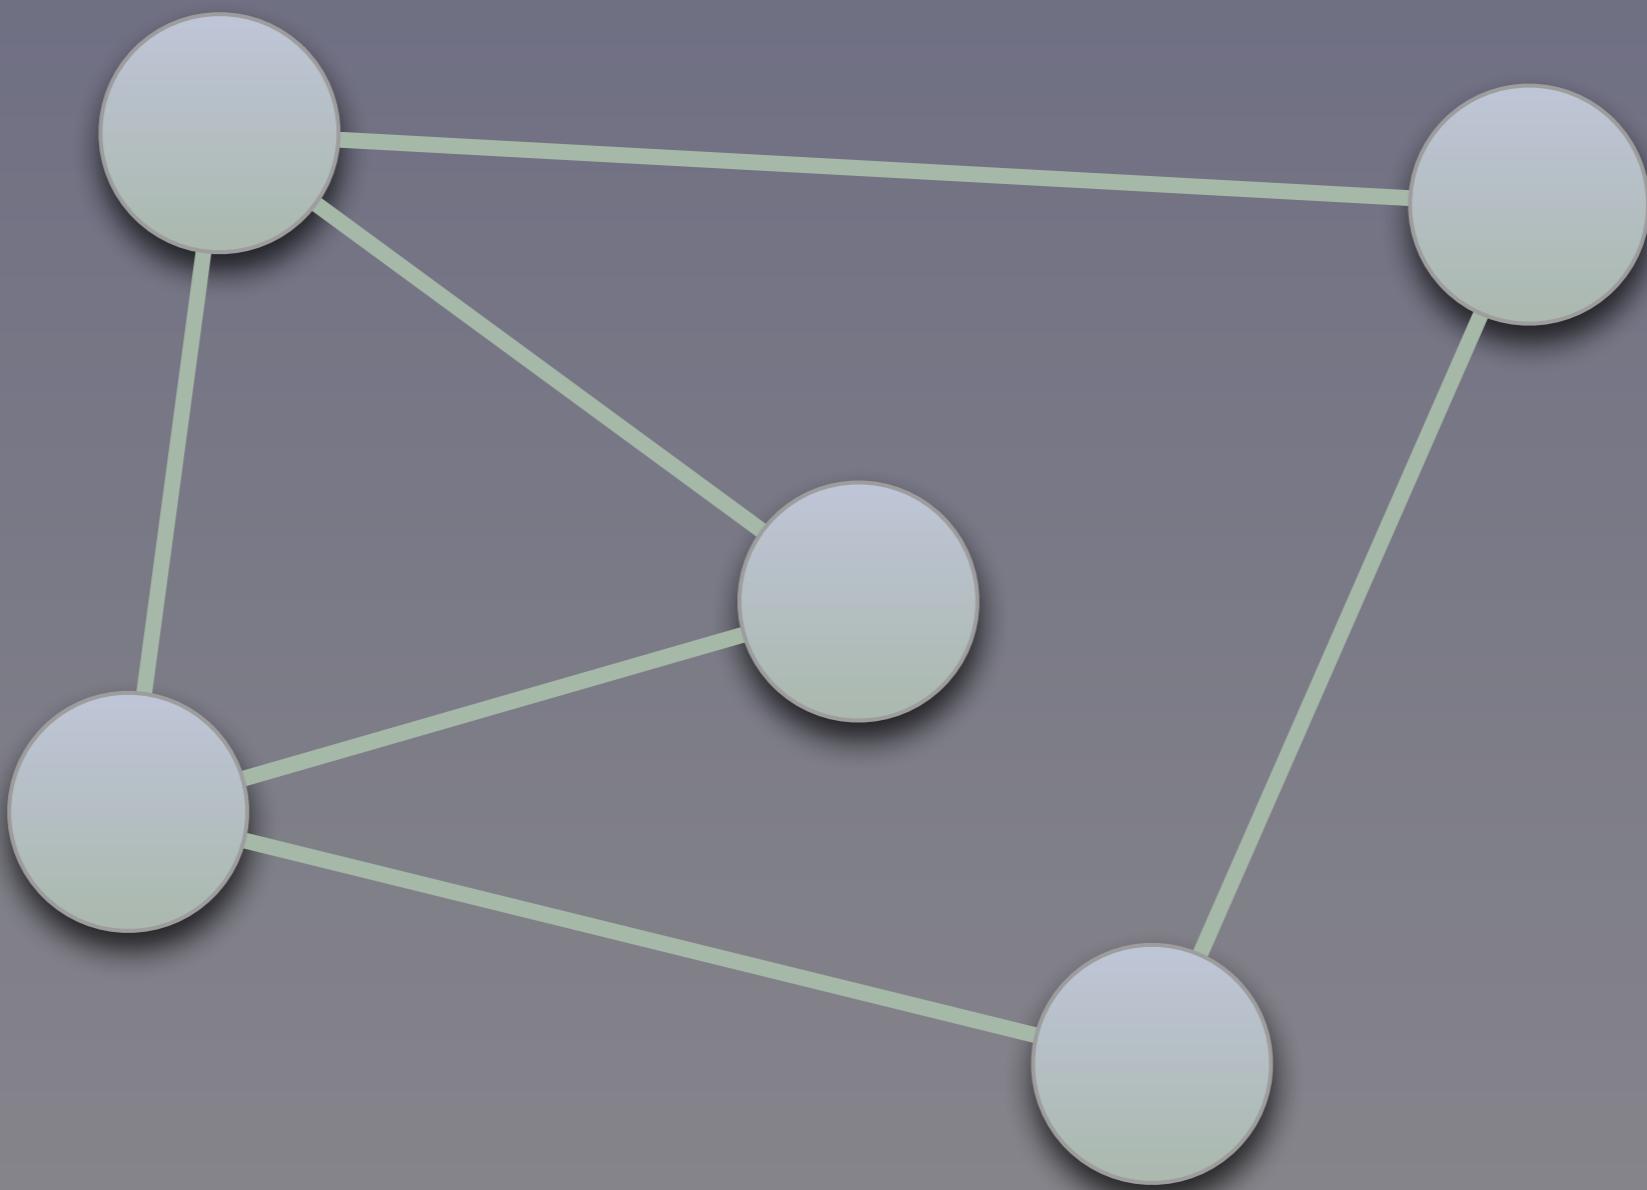
Pregel is inspired by
*Bulk Synchronous
Parallel (BSP)*.



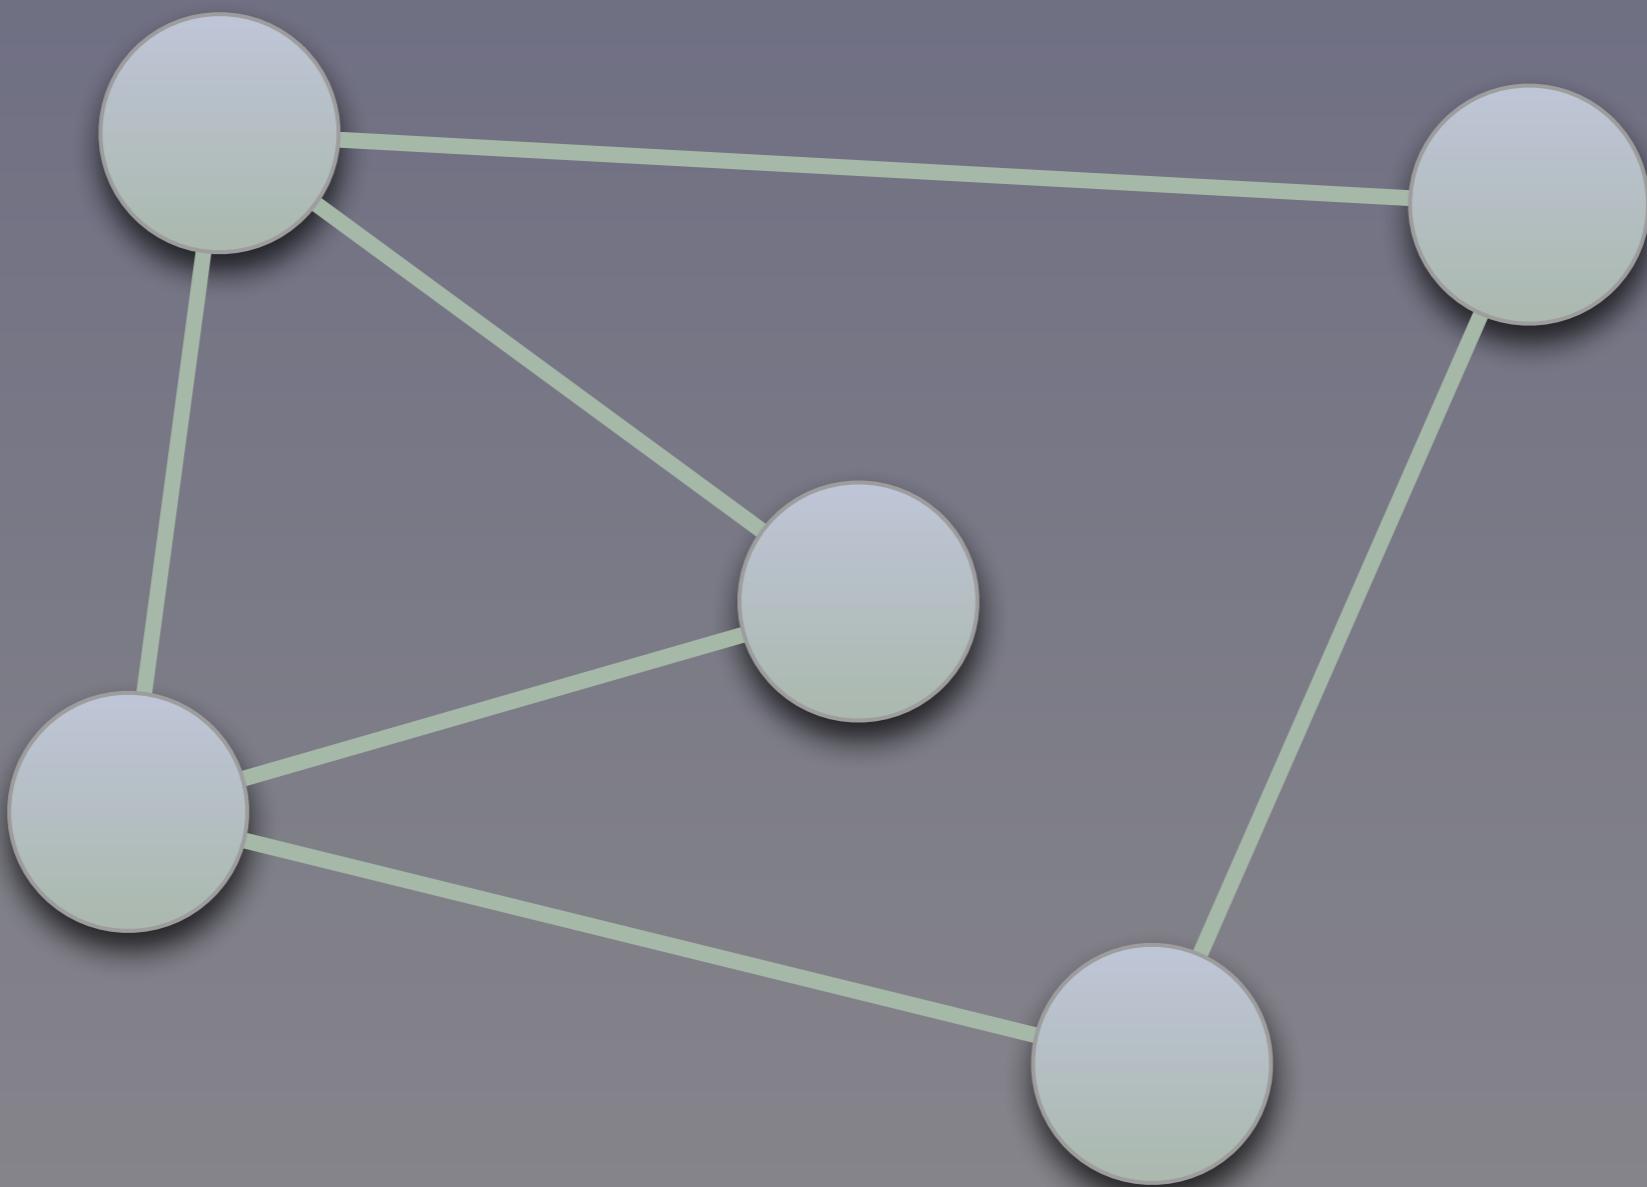
At time t_i , each node can:
receive messages,

Tuesday, June 19, 12

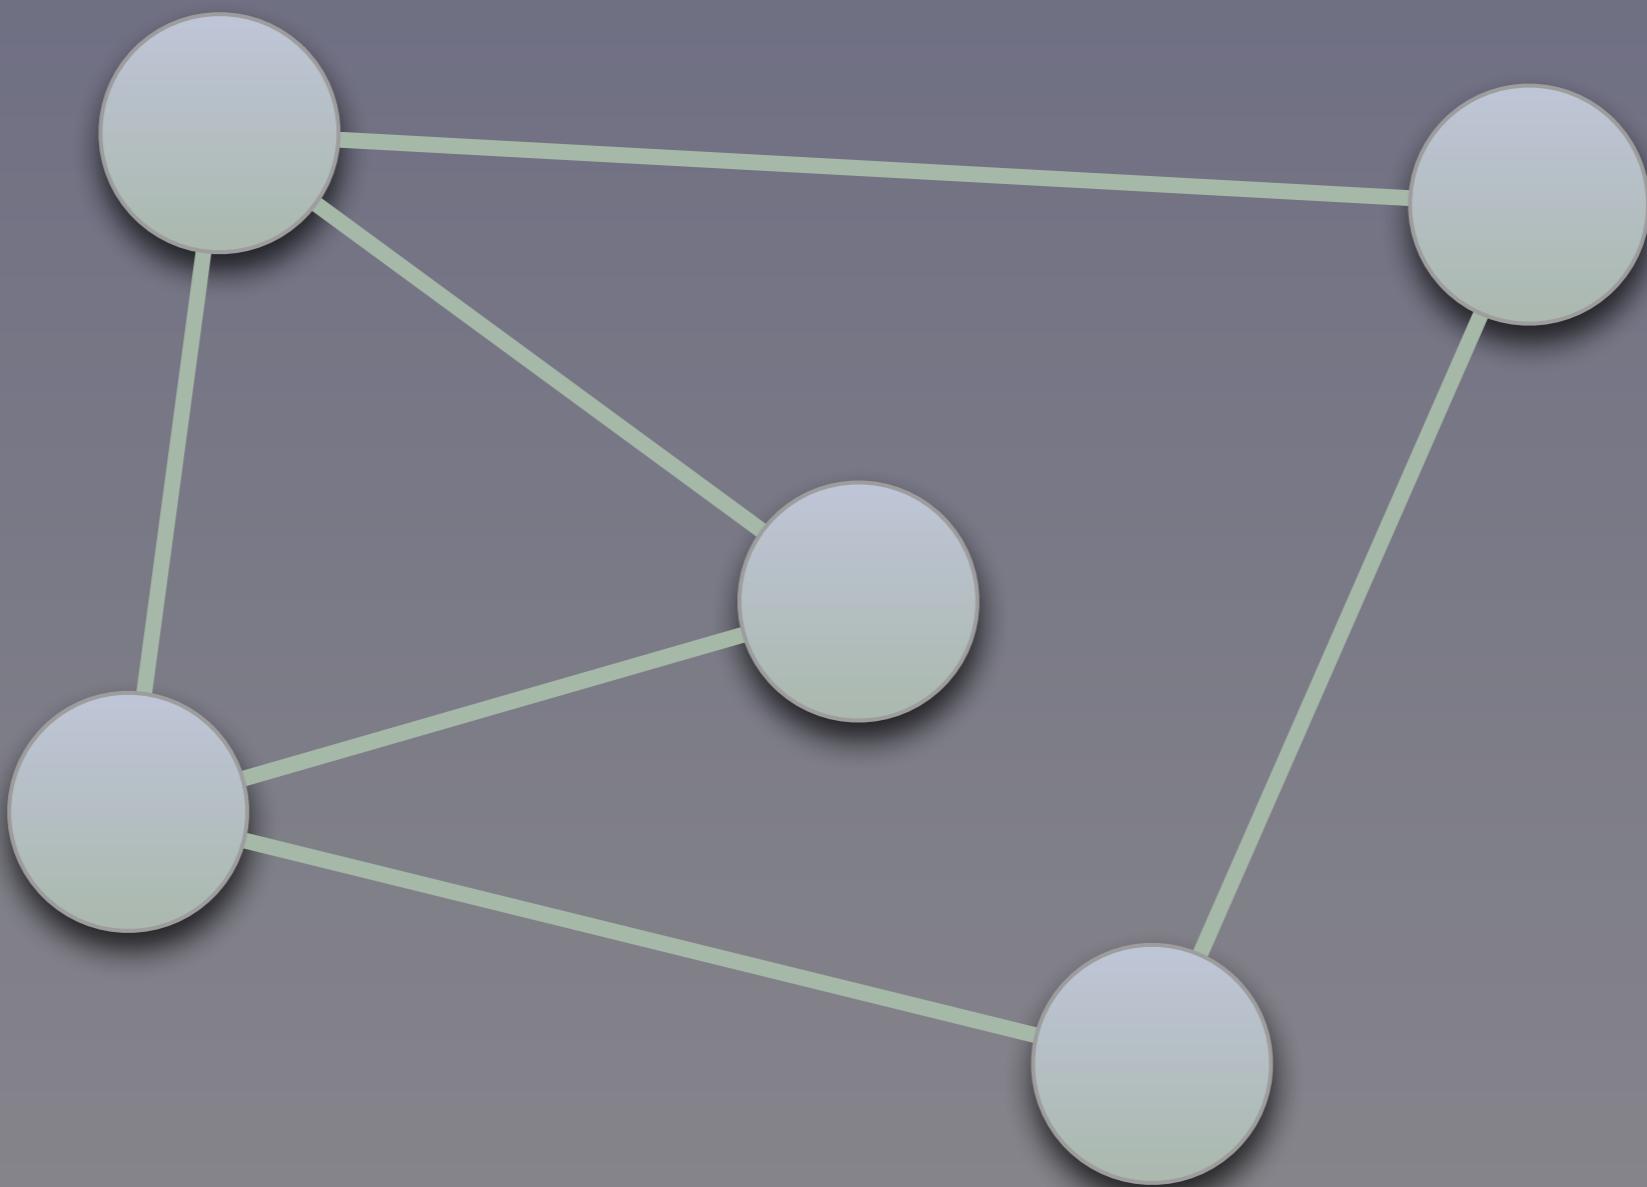
The system marches in synchronous steps, where processing happens at each time “ t_i ”.



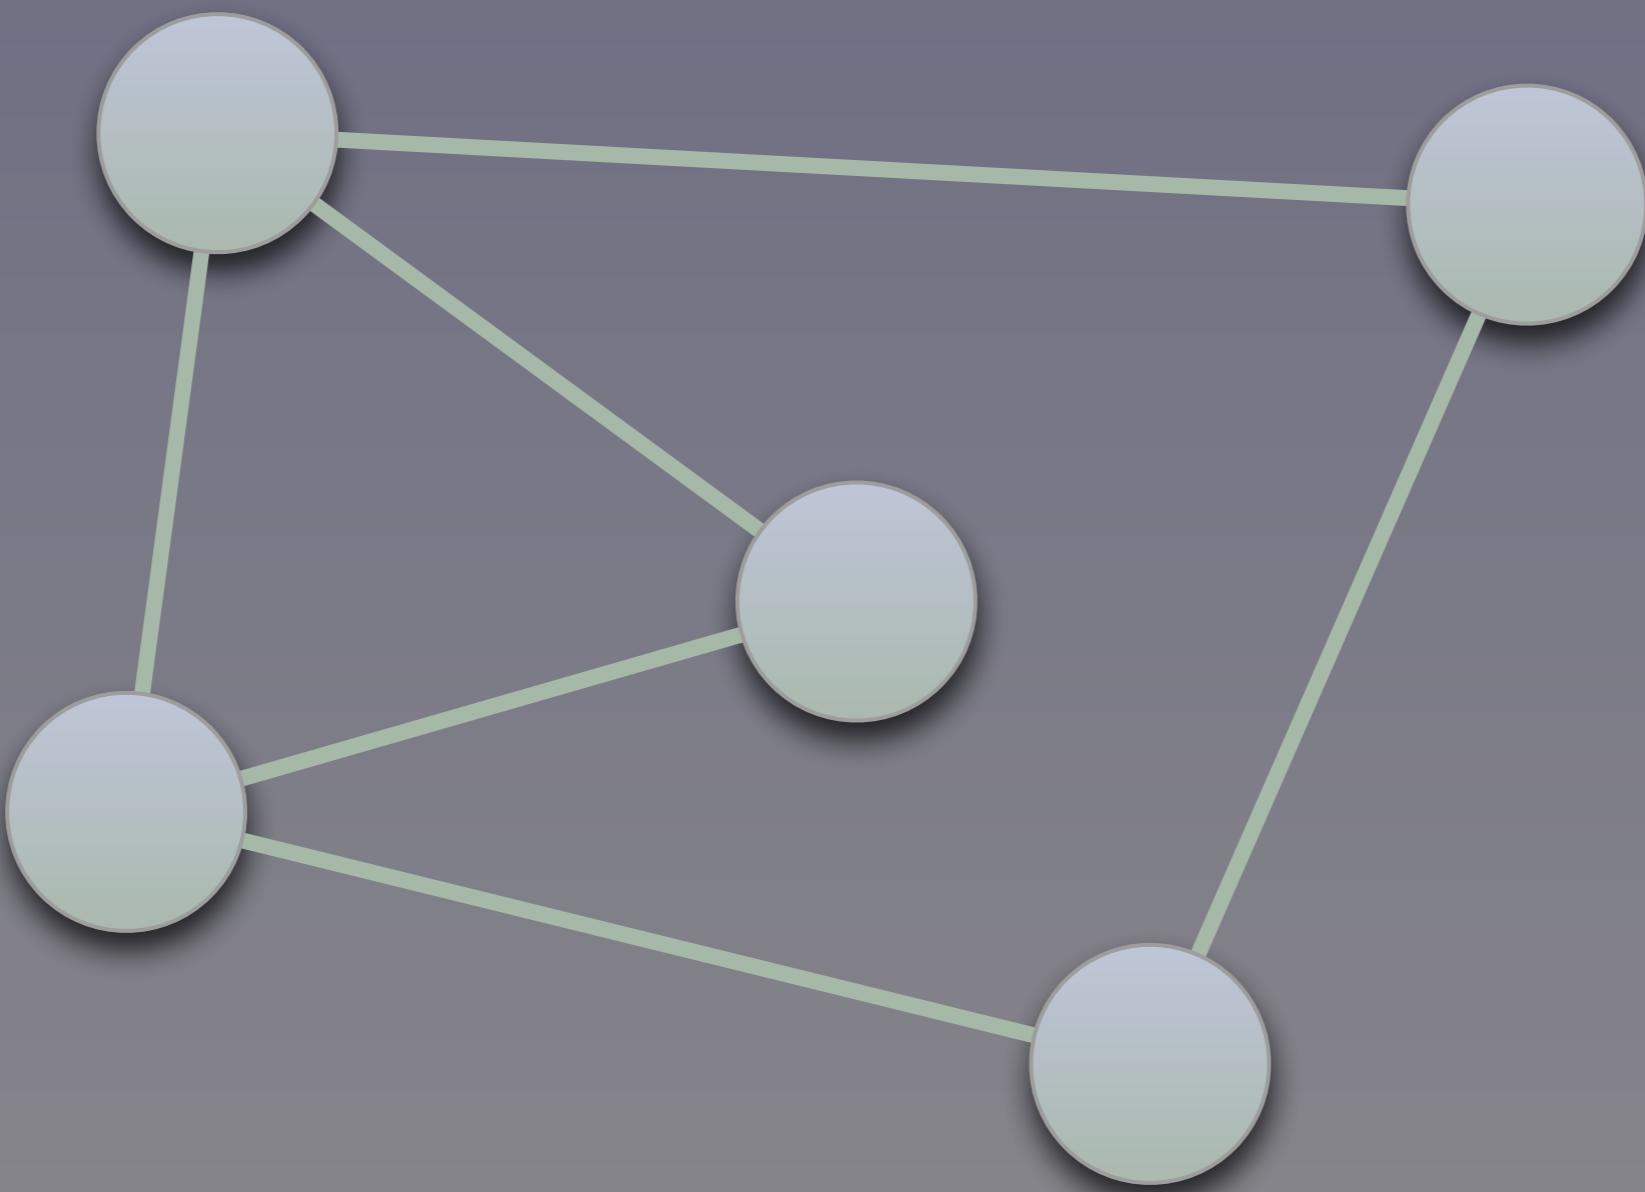
At time t_i , each node can:
send messages,



At time t_i , each node can:
modify its state,



At time t_i , each node can:
modify the edges' states,



At time t_i , each node can:
mutate the graph topology.



Information

[Overview](#)
[Downloads](#)
[Who We Are](#)
[Privacy Policy](#)

Resources

[Developers](#)
[Mailing Lists](#)
[Issue Tracking](#)
[IRC channel](#)

Documentation

[Getting Started](#)
[Hama on Clouds](#)
[Hama BSP tutorial](#)
[Benchmarks](#)
[Presentations](#)
[Research Papers](#)
[Books](#)

Related Projects

[Hadoop](#)

Apache Hama

Apache Hama is a pure BSP(Bulk Synchronous Parallel) computing framework on top of HDFS (Hadoop Distributed File System) for massive scientific computations such as matrix, graph and network algorithms. Currently, it has the following features:

- Job submission and management interface.
- Multiple tasks per node.
- Input/Output Formatter.
- Checkpoint recovery.
- Support to run in the Cloud using [Apache Whirr](#).
- Support to run with Hadoop YARN.



Apache Giraph



Giraph

[About](#)
[Wiki](#)

Project Information

[Summary](#)
[Team](#)
[Mailing Lists](#)
[License](#)
[Issue Tracking](#)
[Source Repository](#)
[Dependencies](#)
[Reports](#)

Welcome To Apache Incubator Giraph

Web and online social graphs have been rapidly growing in size and scale during the past decade. In 2008, Google estimated that the number of web pages reached over a trillion. Online social networking and email sites, including Yahoo!, Google, Microsoft, Facebook, LinkedIn, and Twitter, have hundreds of millions of users and are expected to grow much more in the future. Processing these graphs plays a big role in relevant and personalized information for users, such as results from a search engine or news in an online social networking site.

Tuesday, June 19, 12

Two BSP projects on Apache:

<http://incubator.apache.org/hama/>

<http://incubator.apache.org/giraph/>

Graph Databases

- Neo4J
- Allegro Graph
- FlockDB
- Titan
- GraphDB
- ...

A scenic mountain landscape featuring a large, rugged mountain peak in the background, its slopes covered with patches of snow and dark rock. In the middle ground, a dense forest of green coniferous trees lines a hillside. The foreground is a vibrant field of numerous yellow wildflowers, likely sunflowers or a similar species, stretching across the frame.

#4

Event Stream Processing.

Tuesday, June 19, 12

If you want to use Hadoop and process events in “real-time” (for some definition thereof), you’re best option is HBase, the Hadoop-oriented database. You could choose another NoSQL database too, depending on how much Hadoop interop you want. Hadoop’s MapReduce + HDFS itself are batch oriented and unacceptable. You might also want a dedicated event processing system...



Storm

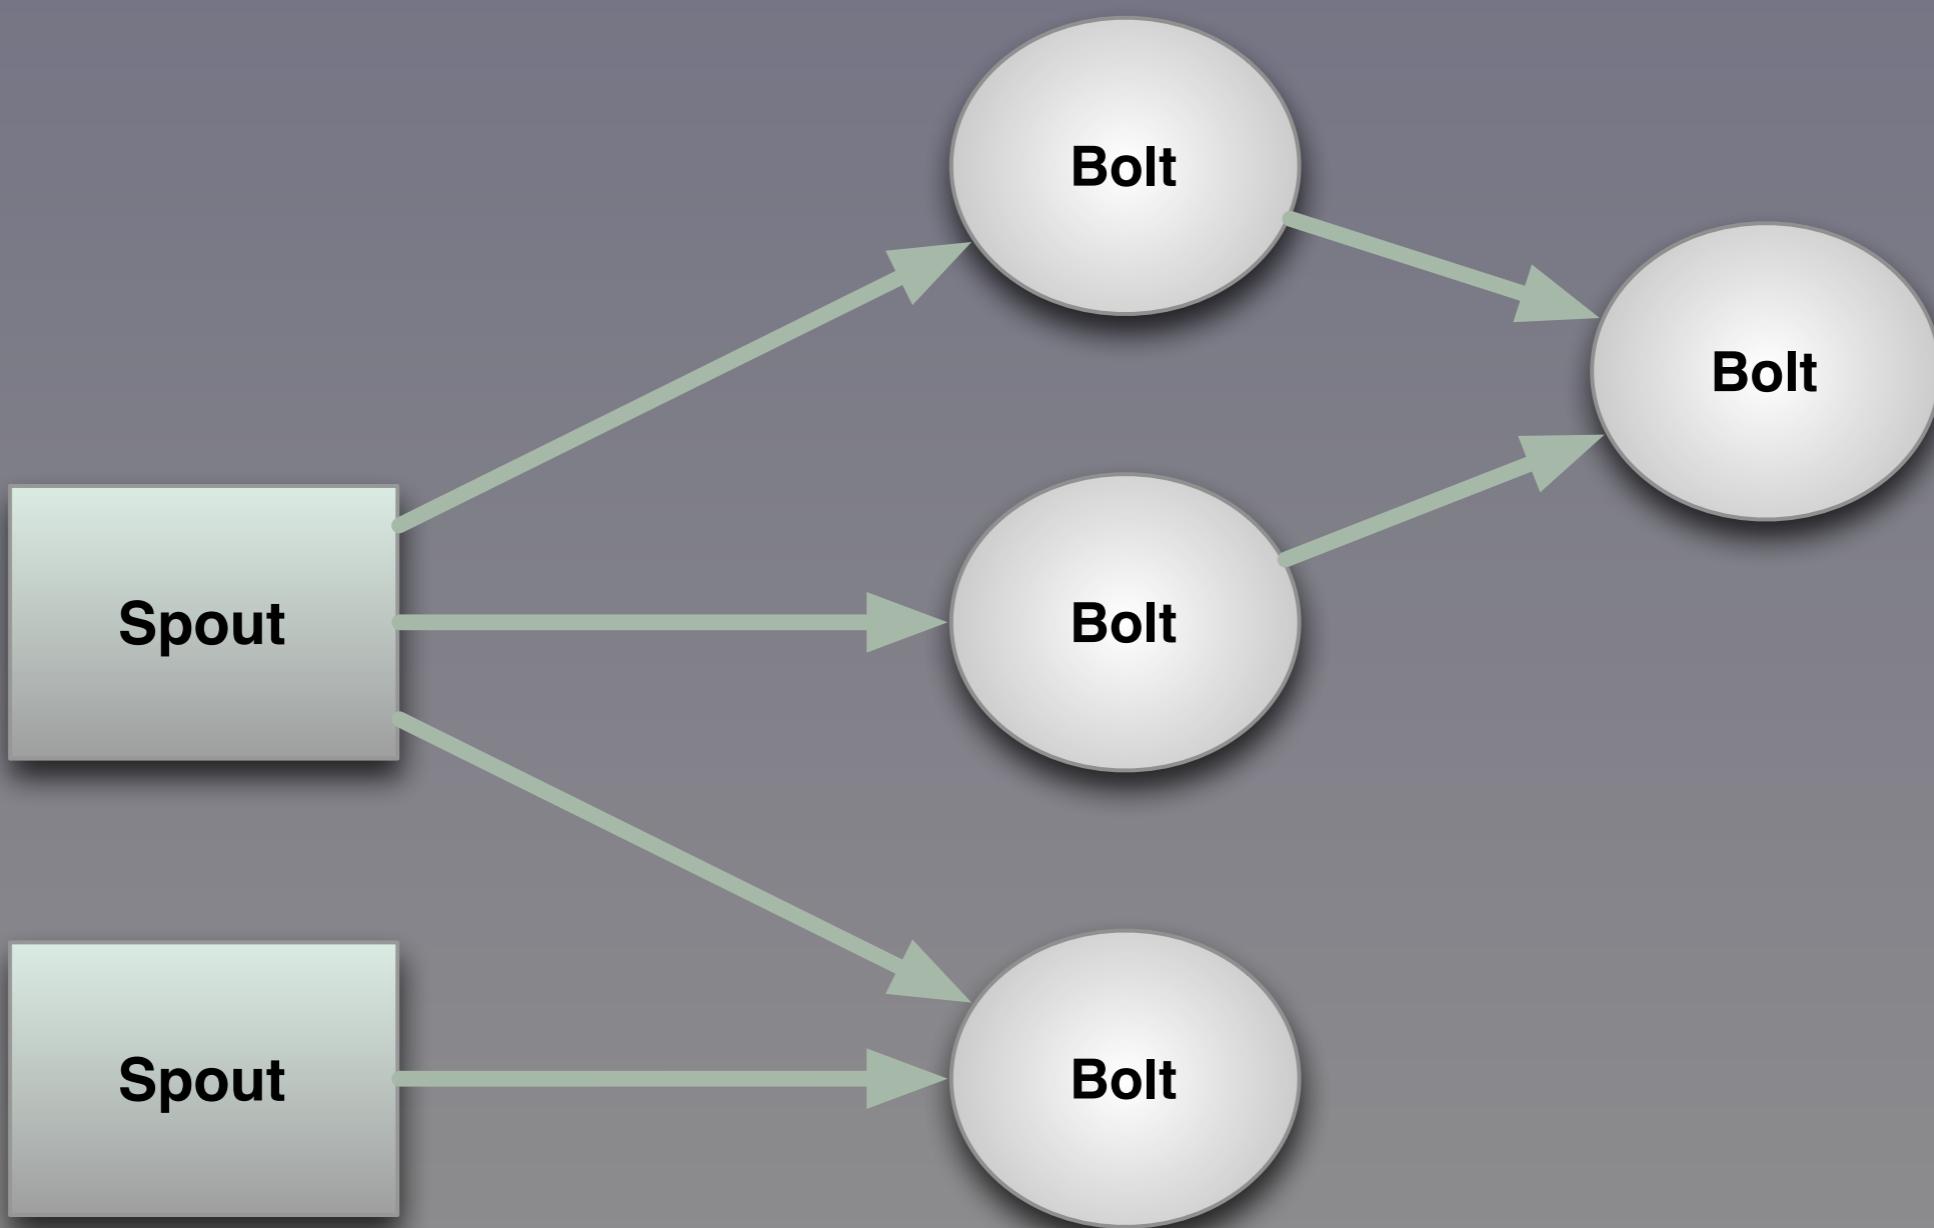
Distributed and fault-tolerant realtime computation



Nathan Marz
Twitter

Tuesday, June 19, 12

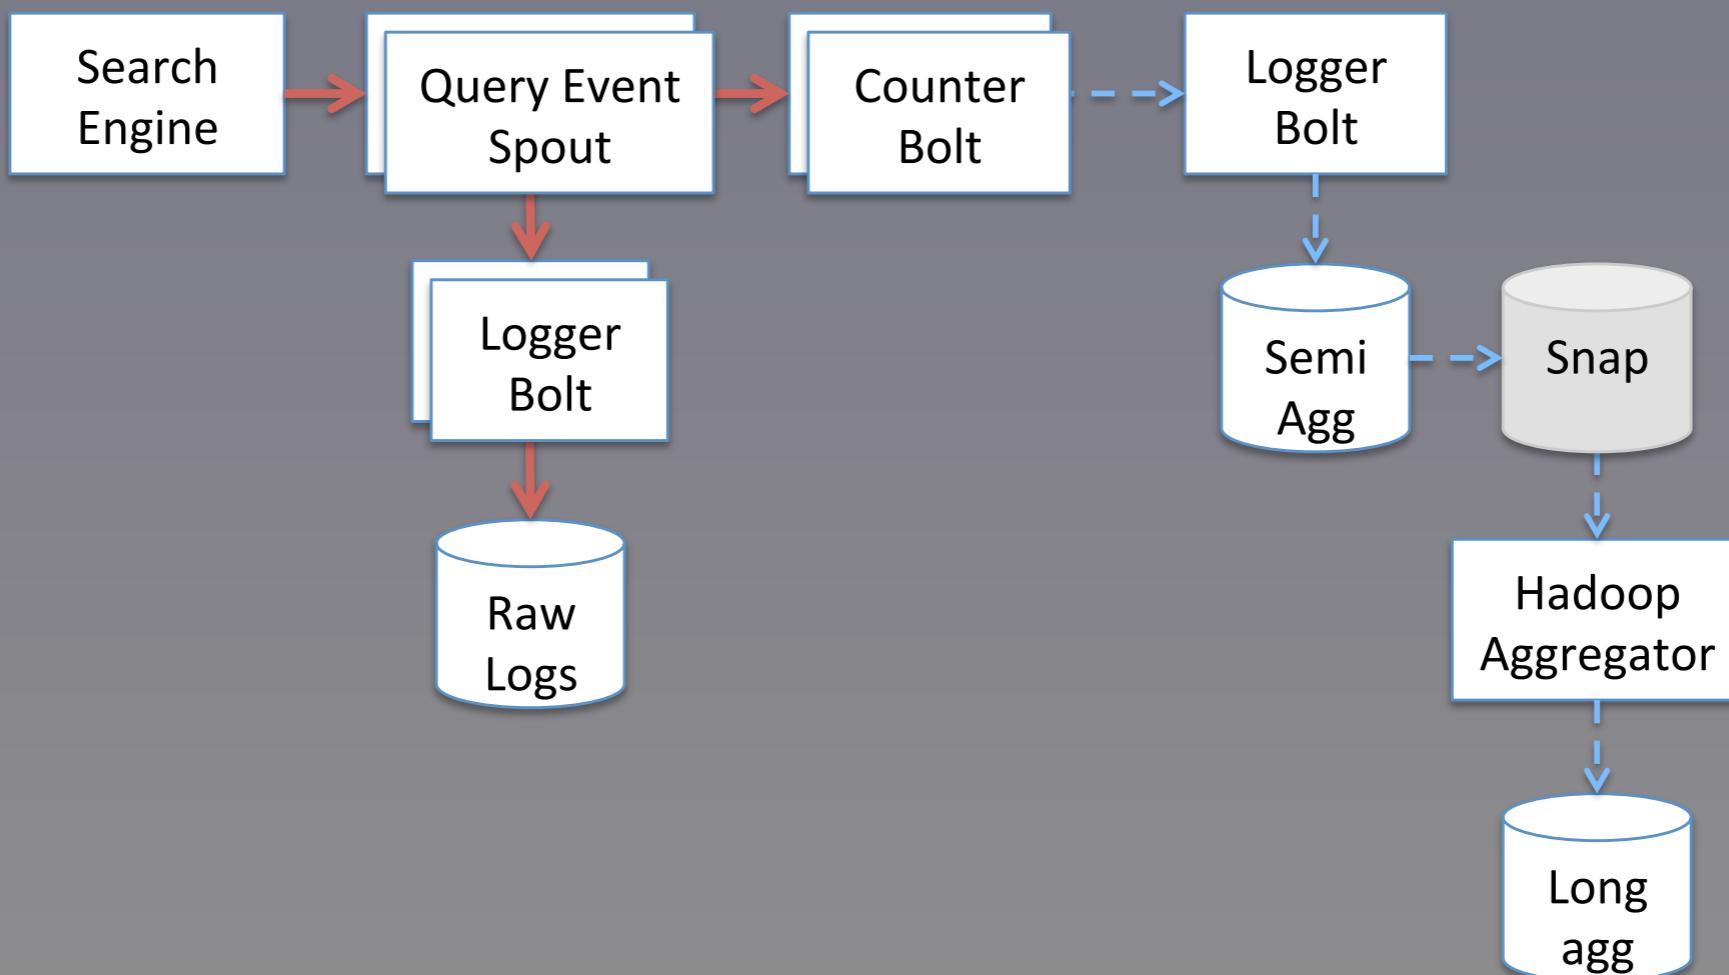
Taken from <https://github.com/strangeloop/2011-slides/raw/master/Marz-Storm.pdf>
See the project page, too: <http://storm-project.net/>



Tuesday, June 19, 12

If you want to use Hadoop and process events in “real-time” (for some definition thereof), you’re best option is HBase, the Hadoop-oriented database. You could choose another NoSQL database too, depending on how much Hadoop interop you want. Hadoop’s MapReduce + HDFS itself are batch oriented and unacceptable. You might also want a dedicated event processing system...

Storm + Hadoop



Tuesday, June 19, 12

You can combine Storm with Hadoop to get the best of both worlds.



A Manifesto...

Tuesday, June 19, 12

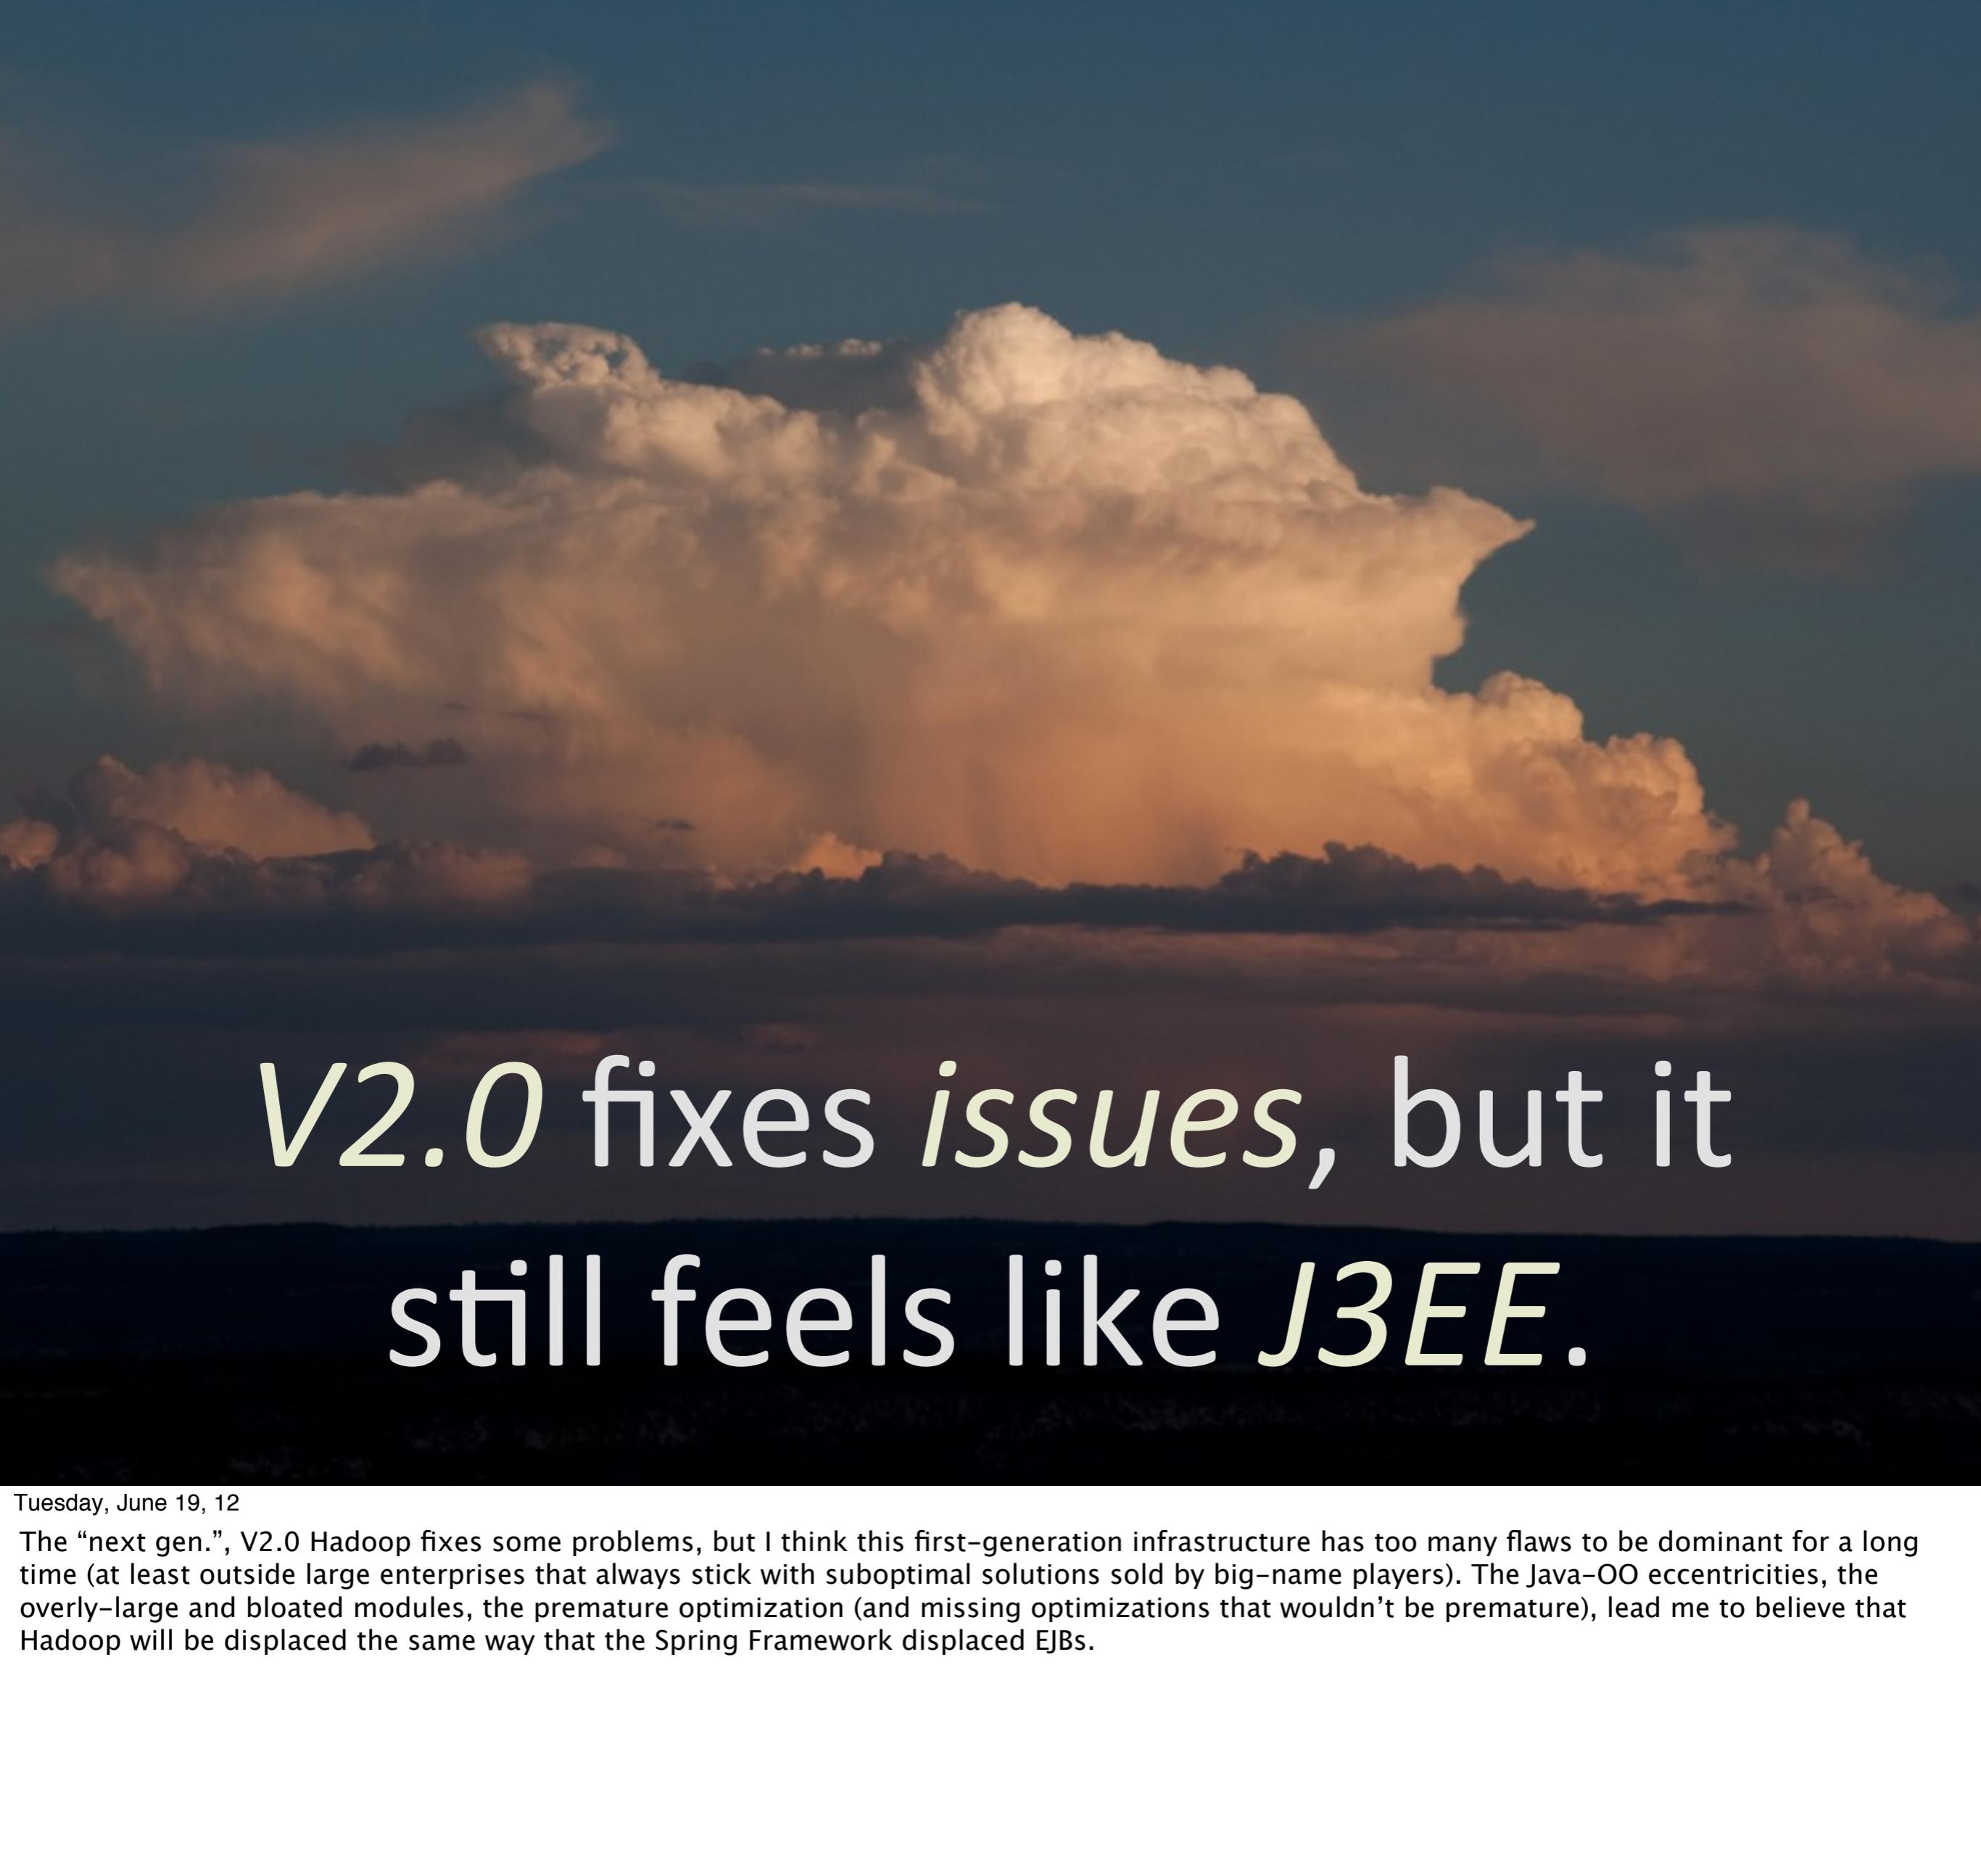
So, I think we have an opportunity...



Hadoop is the Enterprise Java Beans of our time.

Tuesday, June 19, 12

I worked with EJBs a decade ago. Like EJBs, Hadoop has an invasive API that obscures your business logic and reusability. There were too many configuration options in XML files. The framework “paradigm” is a poor fit for most problems (like soft real time systems and most algorithms beyond Word Count). Internally, EJB implementations were inefficient and hard to optimize, because they relied on poorly considered object boundaries that muddled more natural boundaries and created large-scale, monolithic modules with few abstractions for extension and optimization points. I’ve also argued in other presentations and my “FP for Java Devs” book that OOP is a poor modularity tool... The fact is, Hadoop reminds me of EJBs in almost every way. It works okay and people do get stuff done, but just as the Spring Framework brought an essential rethinking to Enterprise Java, I think there is an essential rethink that needs to happen in Big Data. The FP community is well positioned to create the next generation.



V2.0 fixes *issues*, but it
still feels like J3EE.

Tuesday, June 19, 12

The “next gen.”, V2.0 Hadoop fixes some problems, but I think this first-generation infrastructure has too many flaws to be dominant for a long time (at least outside large enterprises that always stick with suboptimal solutions sold by big-name players). The Java–OO eccentricities, the overly-large and bloated modules, the premature optimization (and missing optimizations that wouldn’t be premature), lead me to believe that Hadoop will be displaced the same way that the Spring Framework displaced EJBs.



Stop using Java.

Tuesday, June 19, 12

Seriously, it's the worst possible language for Big Data. It promotes bloat and bad abstractions. Without a REPL, you can't experiment. You are penalizing yourself by using it.
You don't have to rewrite your old code, you can just call it from Scala or Clojure, then use those languages (or pick anything else modern, like JRuby and Jython or a non-JVM, functional language).

Assess your needs.

Use optimal tools.



Tuesday, June 19, 12

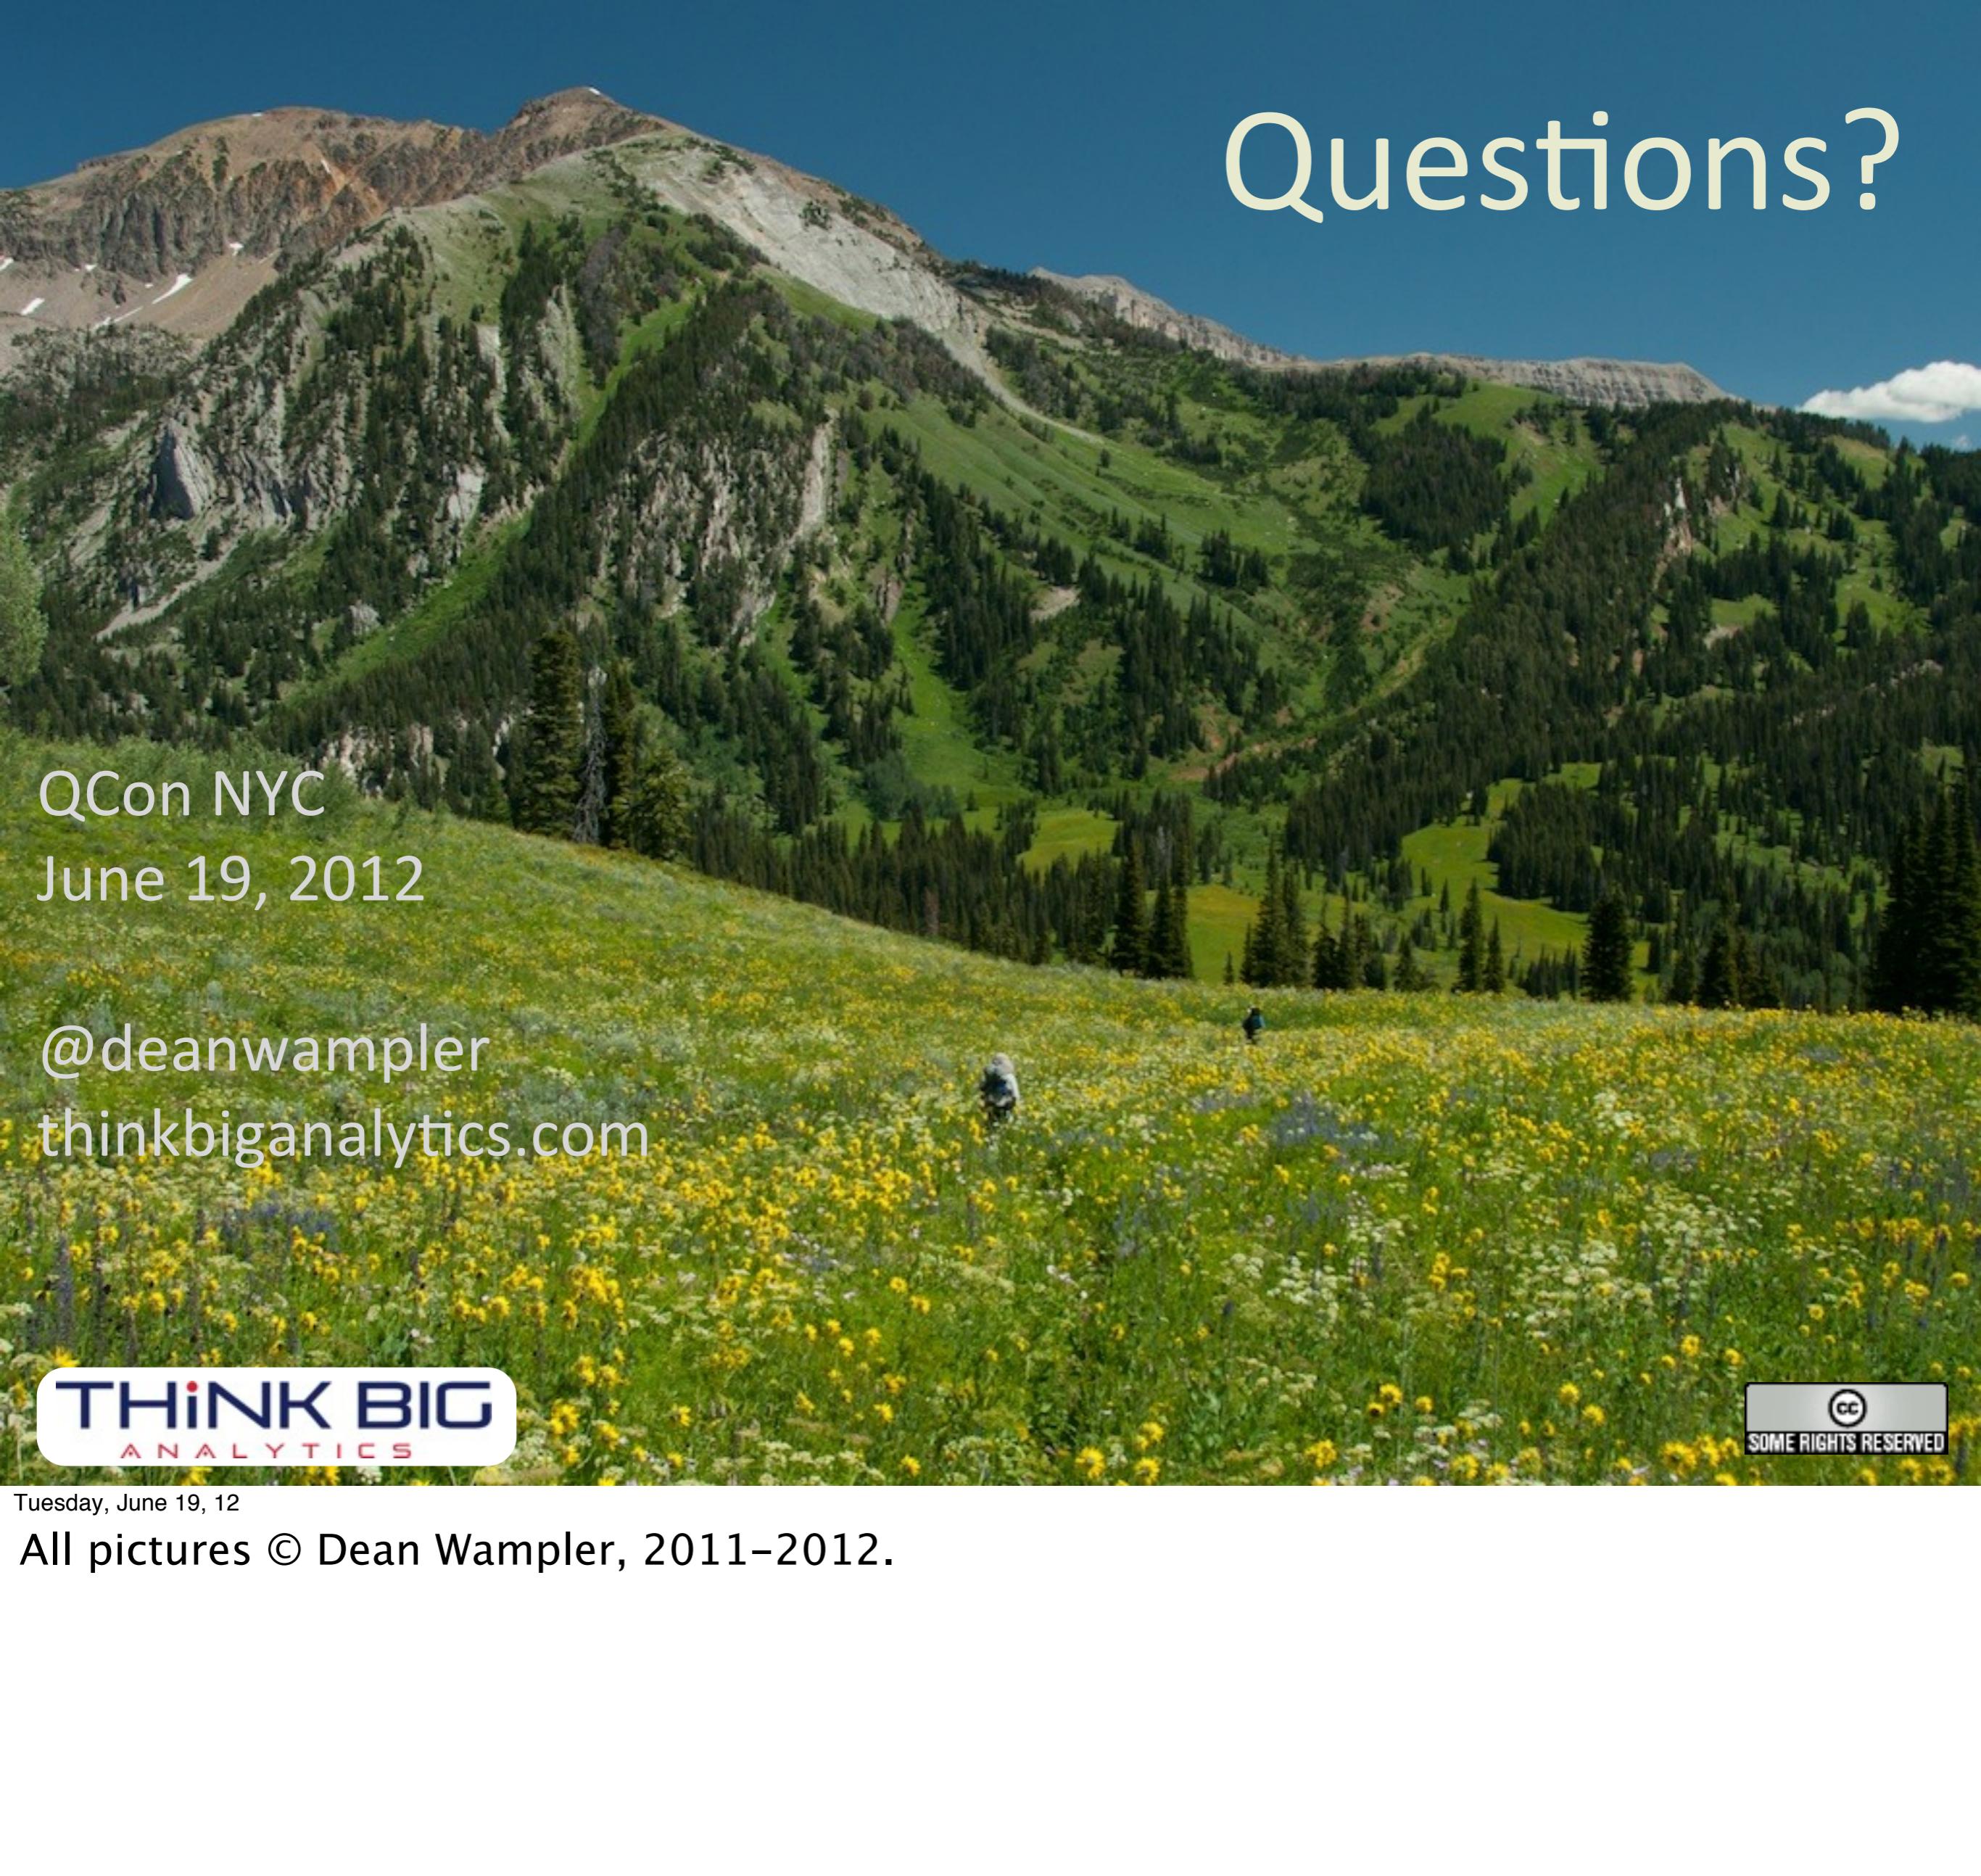
Hadoop and (most) of the ecosystem tools DO have their place, but understand your needs first and then pick the tools that make sense. Skepticism is healthy. If you decide you do need Hadoop scalability (and many people do), then use higher-level languages and tools like Cascading (or better, Cascalog or Scalding) or Hive and Pig. And don't assume you should use every Hadoop ecosystem tool; some are garbage.

Consider Storm and BSP



Tuesday, June 19, 12

Storm provides “real-time”, event-stream processing and BSP (Bulk Synchronous Parallel) is one of several emerging, performant technologies for graph algorithms.

A wide-angle photograph of a mountain range. In the foreground, a hillside is covered in vibrant yellow wildflowers and green grass. Two small figures of people are walking across this field. Behind them, the mountains rise with dense green forests on their slopes. The sky is a clear, bright blue.

Questions?

QCon NYC
June 19, 2012

@deanwampler
thinkbiganalytics.com



Tuesday, June 19, 12

All pictures © Dean Wampler, 2011-2012.