

Ray - Scalability from a Laptop to a Cluster

Dean Wampler - Nov 6, 2020

dean@deanwampler.com

[@deanwampler](https://twitter.com/deanwampler)

ray.io

[Domino Data Lab](#)



System-of-Record for Enterprise Data Science Teams



Accelerate Research

Get self-serve access to the latest tools and scalable compute. Reuse past work and iterate more efficiently.

[Learn More »](#)



Centralize Infrastructure

Manage the availability of powerful data science resources in a secure and governed system-of-record.

[Learn More »](#)



Deploy and Monitor Models

Expedite model consumption with apps, APIs, and more – and ensure their accuracy for key decisions.

[Learn More »](#)



Unify Data Science Teams

Make data science teams more productive and collaborative, and manage their work more efficiently.

[Learn More »](#)





Outline

- Why Ray?
- ML/AI Ray Libraries
- Ray for Microservices
- Adopting Ray





why Ray??





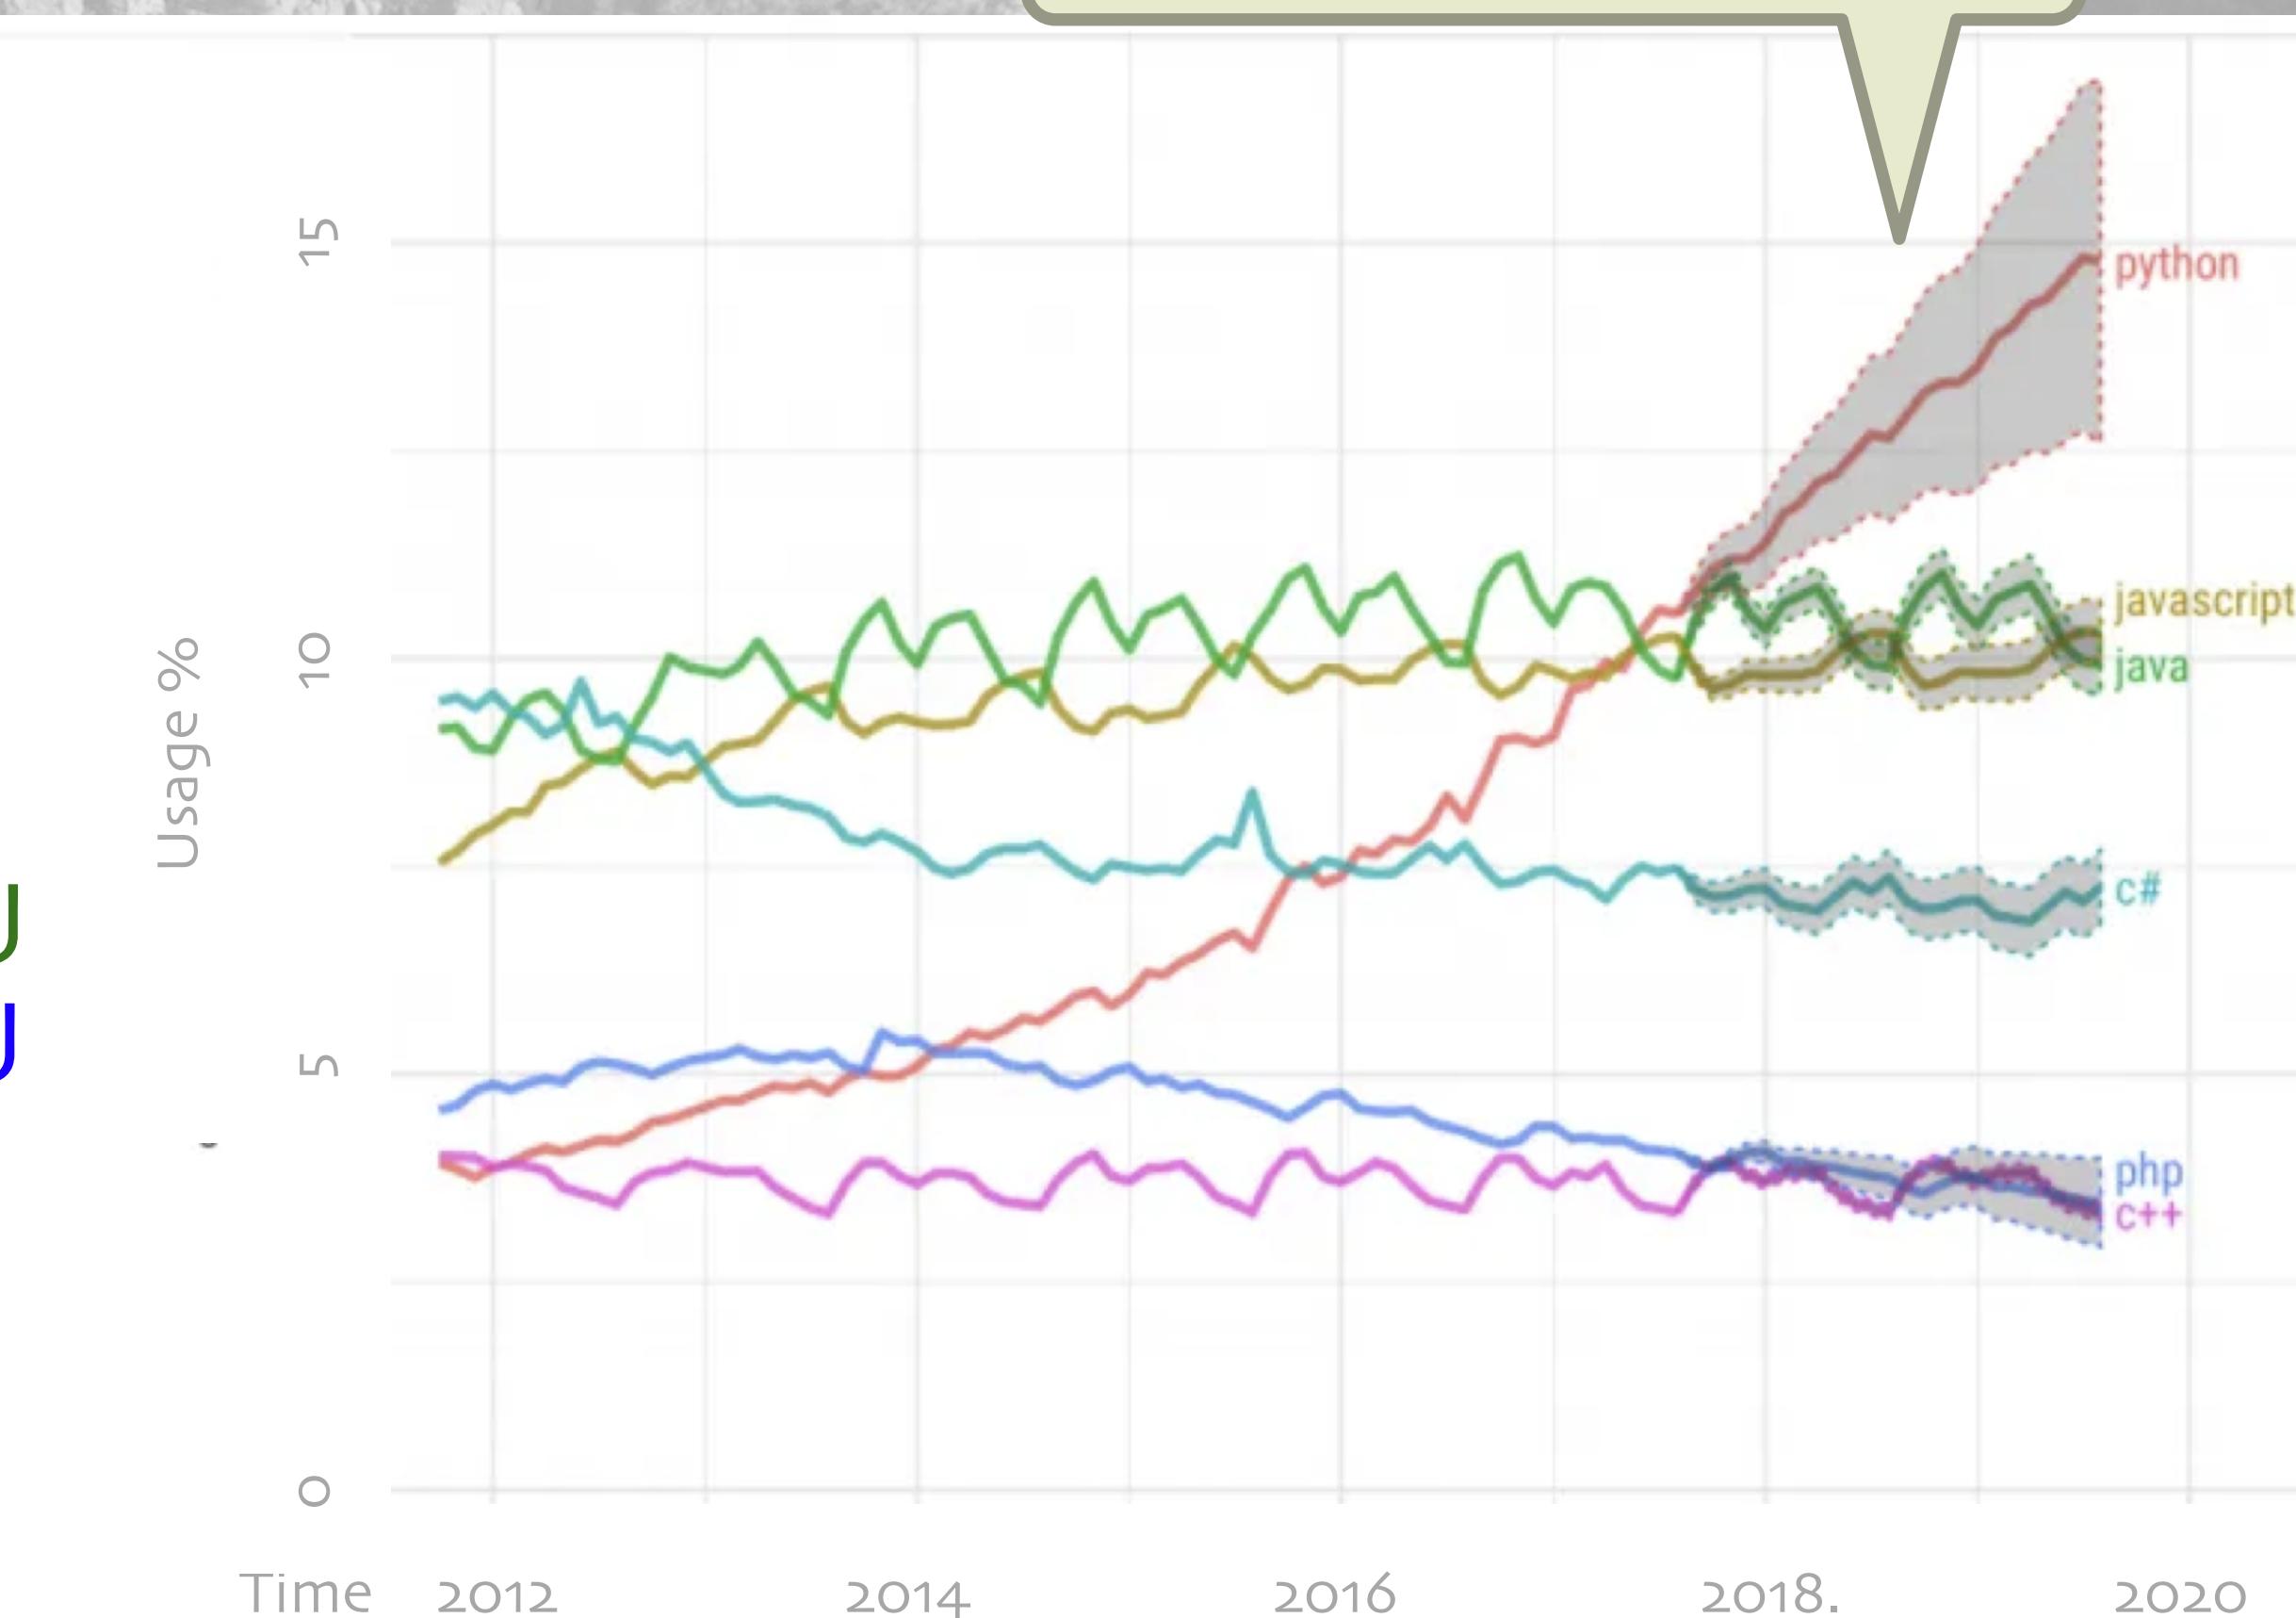
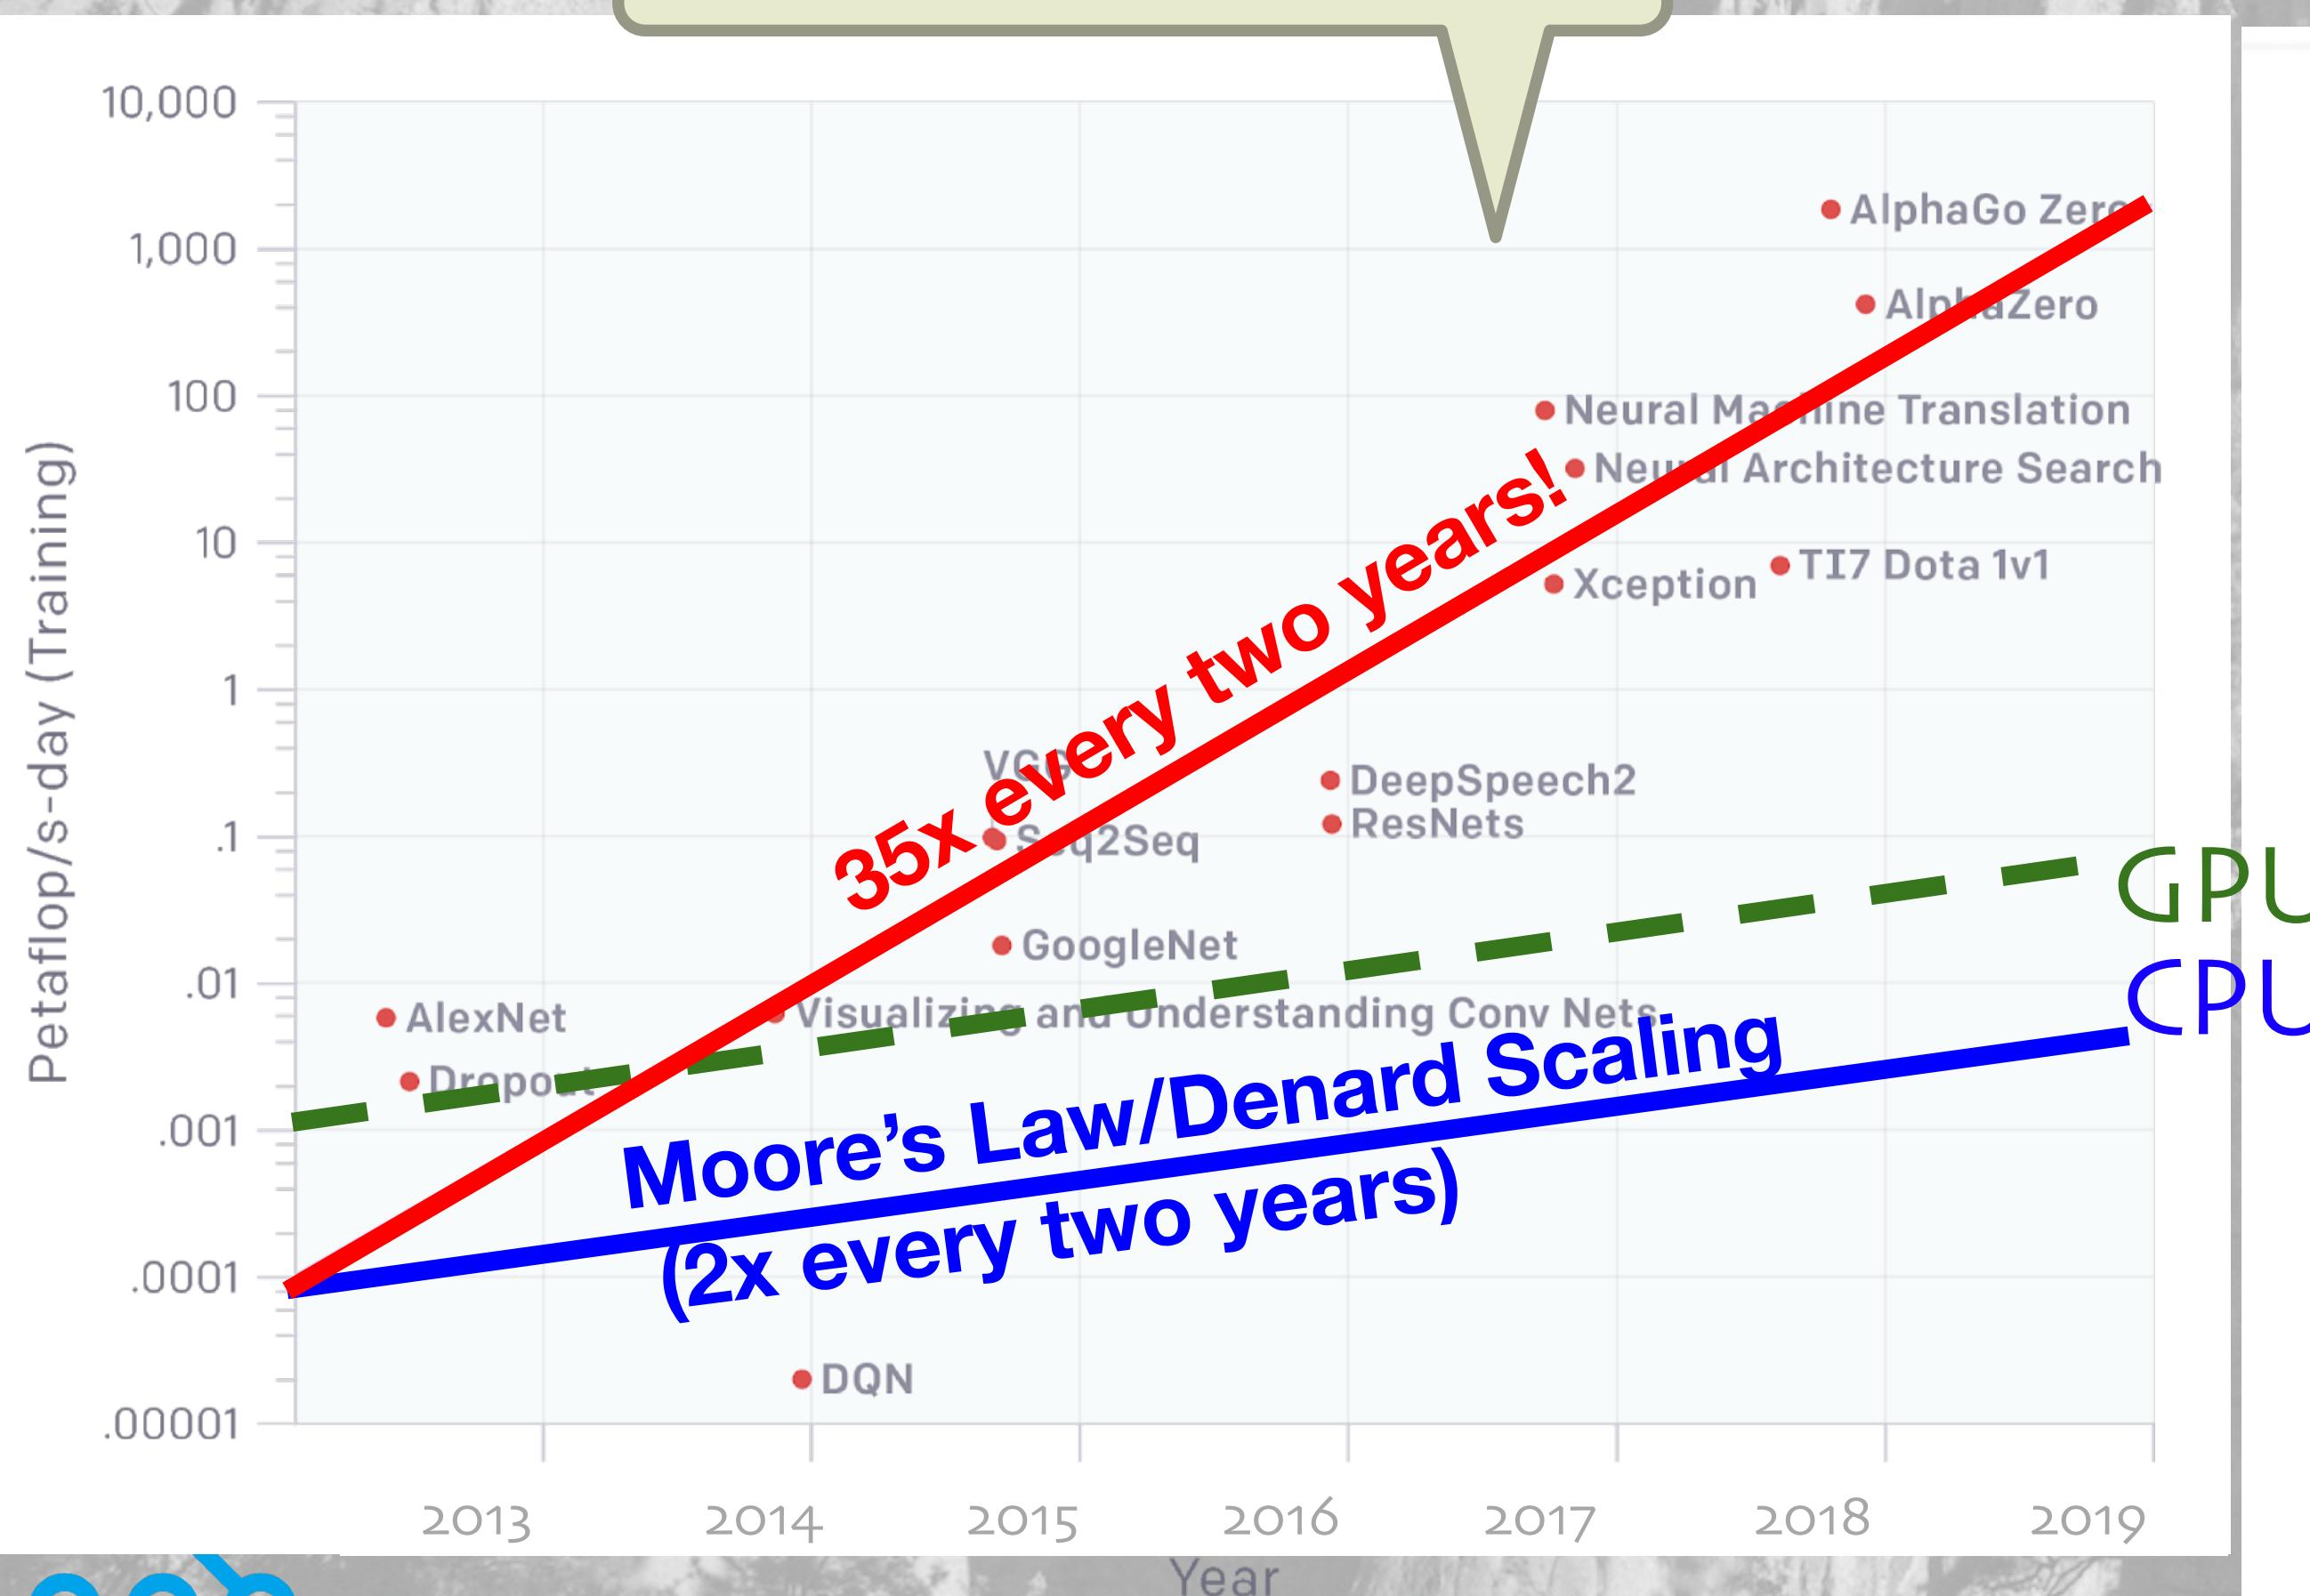
@deanwampler

Two Major Trends

Model sizes and therefore compute requirements outstripping Moore's Law

Hence, there is a pressing need for robust, easy to use solutions for distributed Python

Python growth driven by ML/AI and other data science workloads



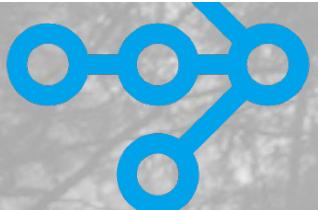
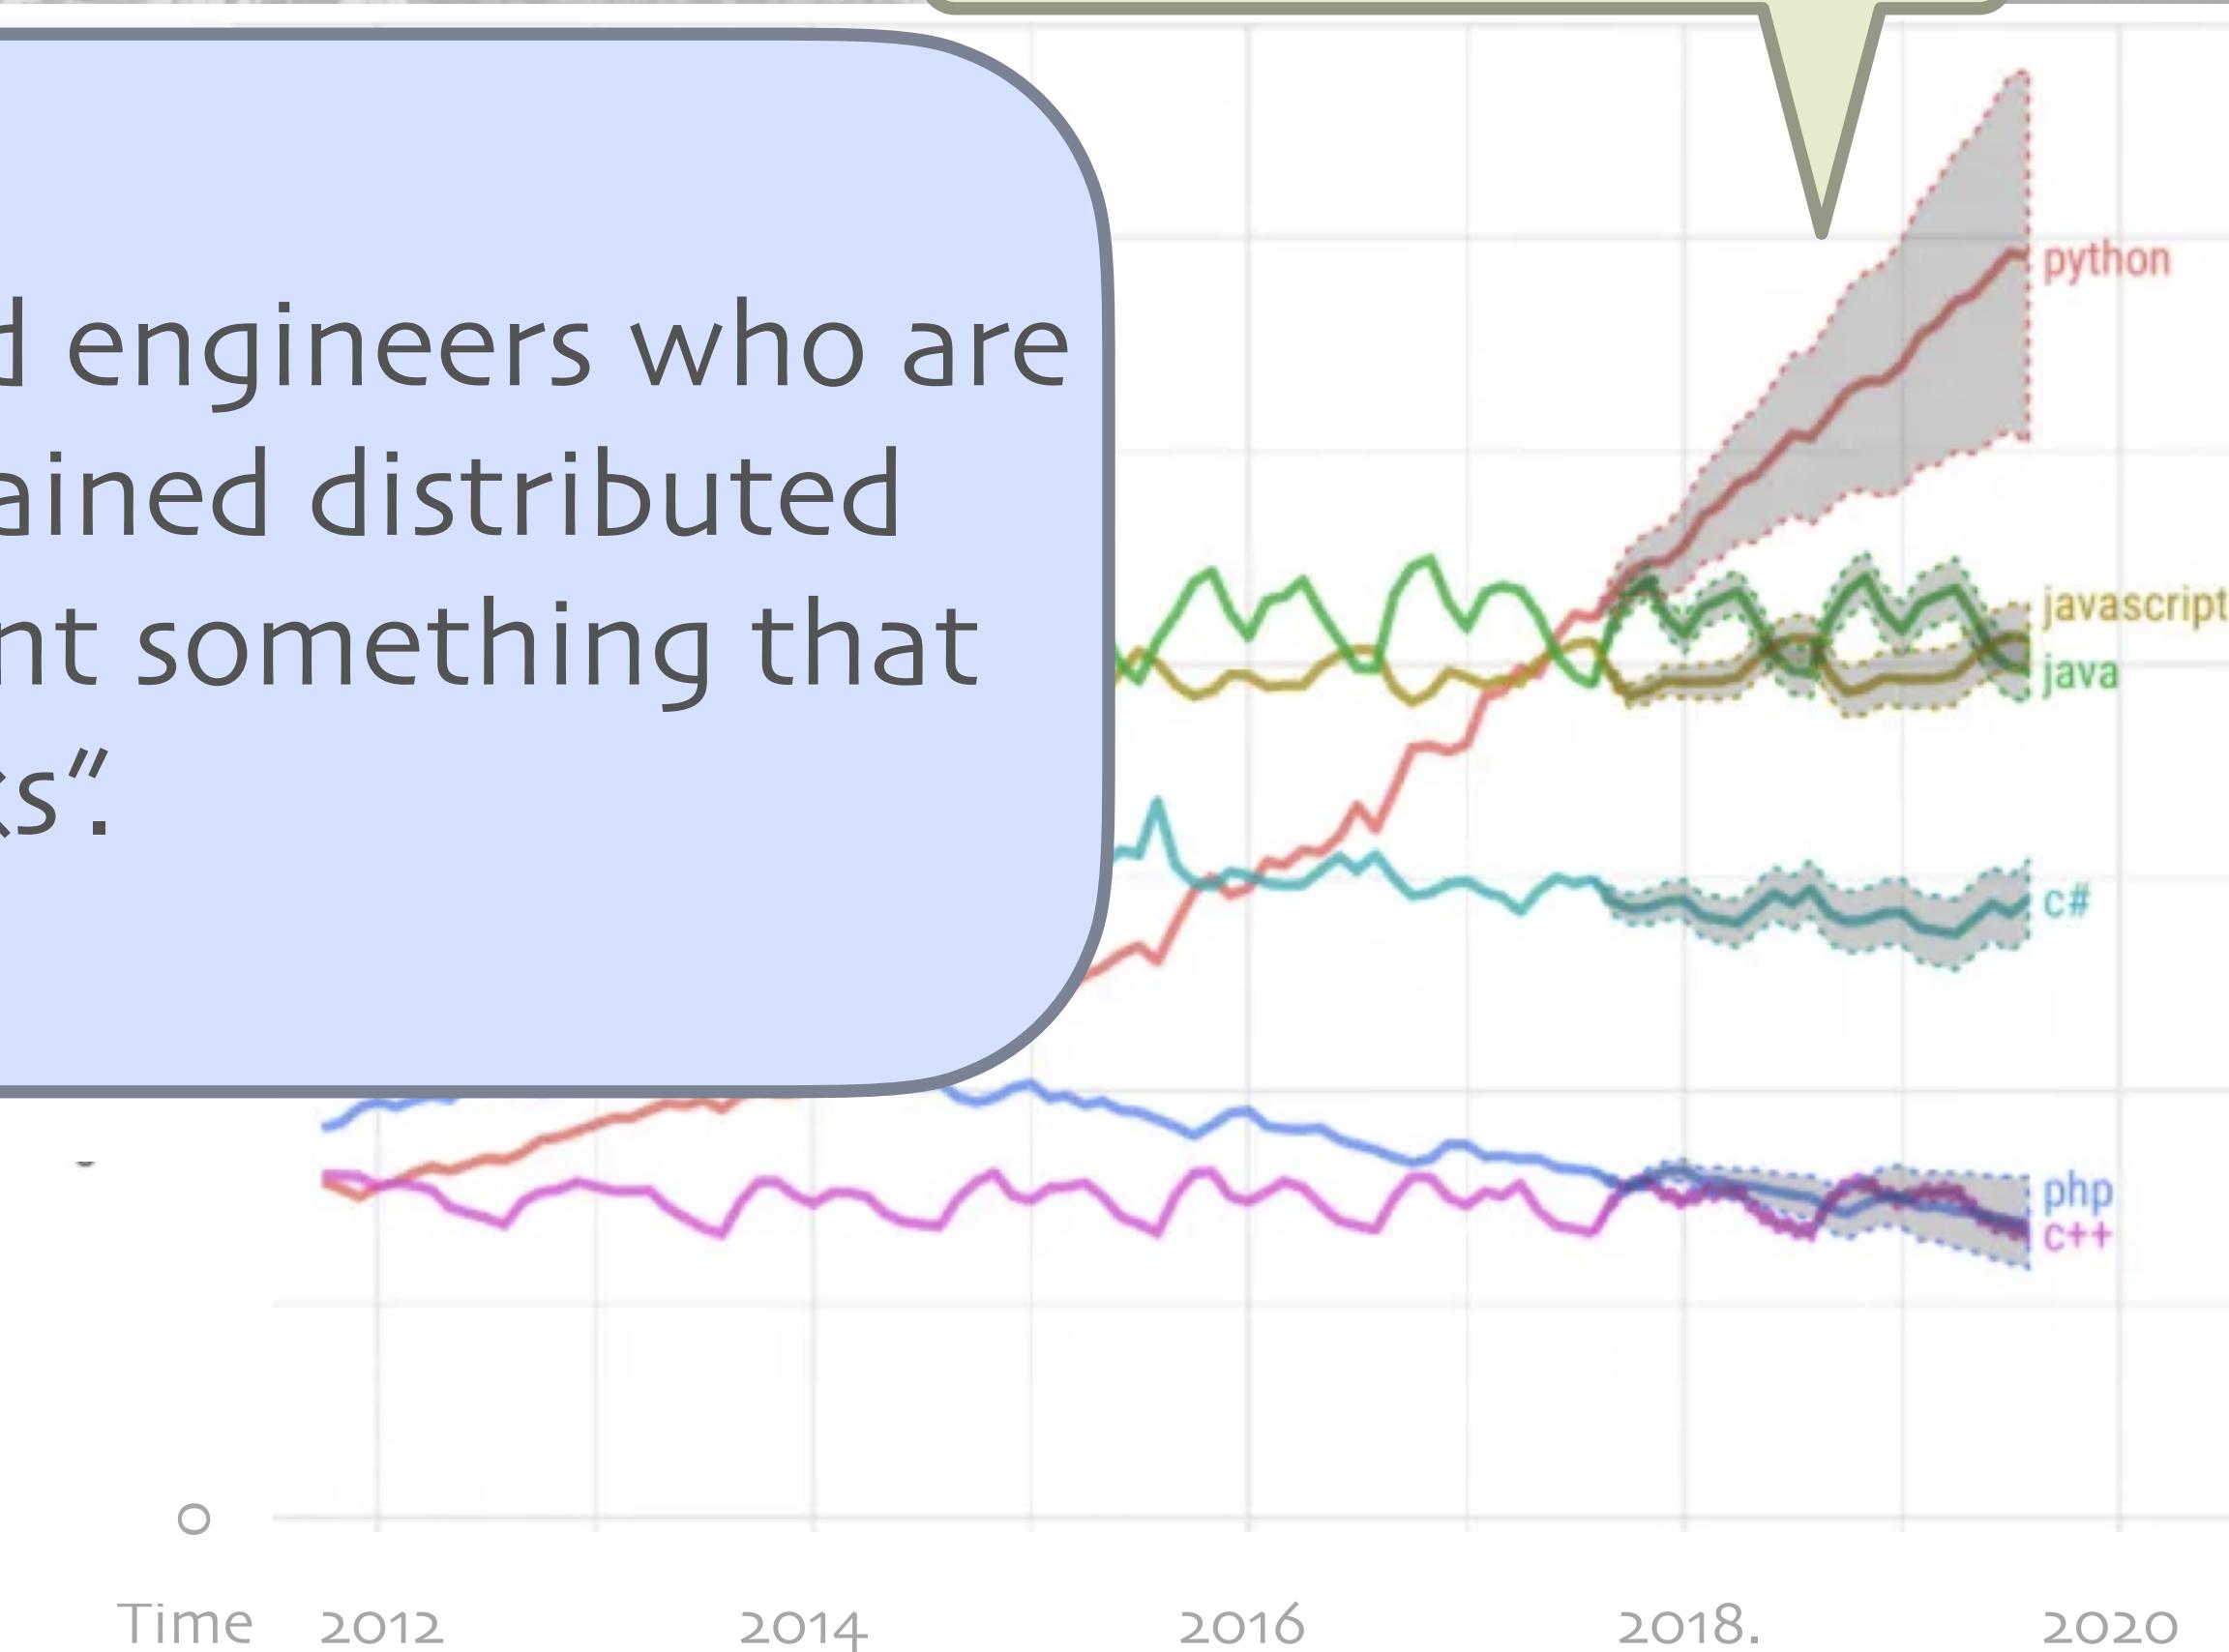
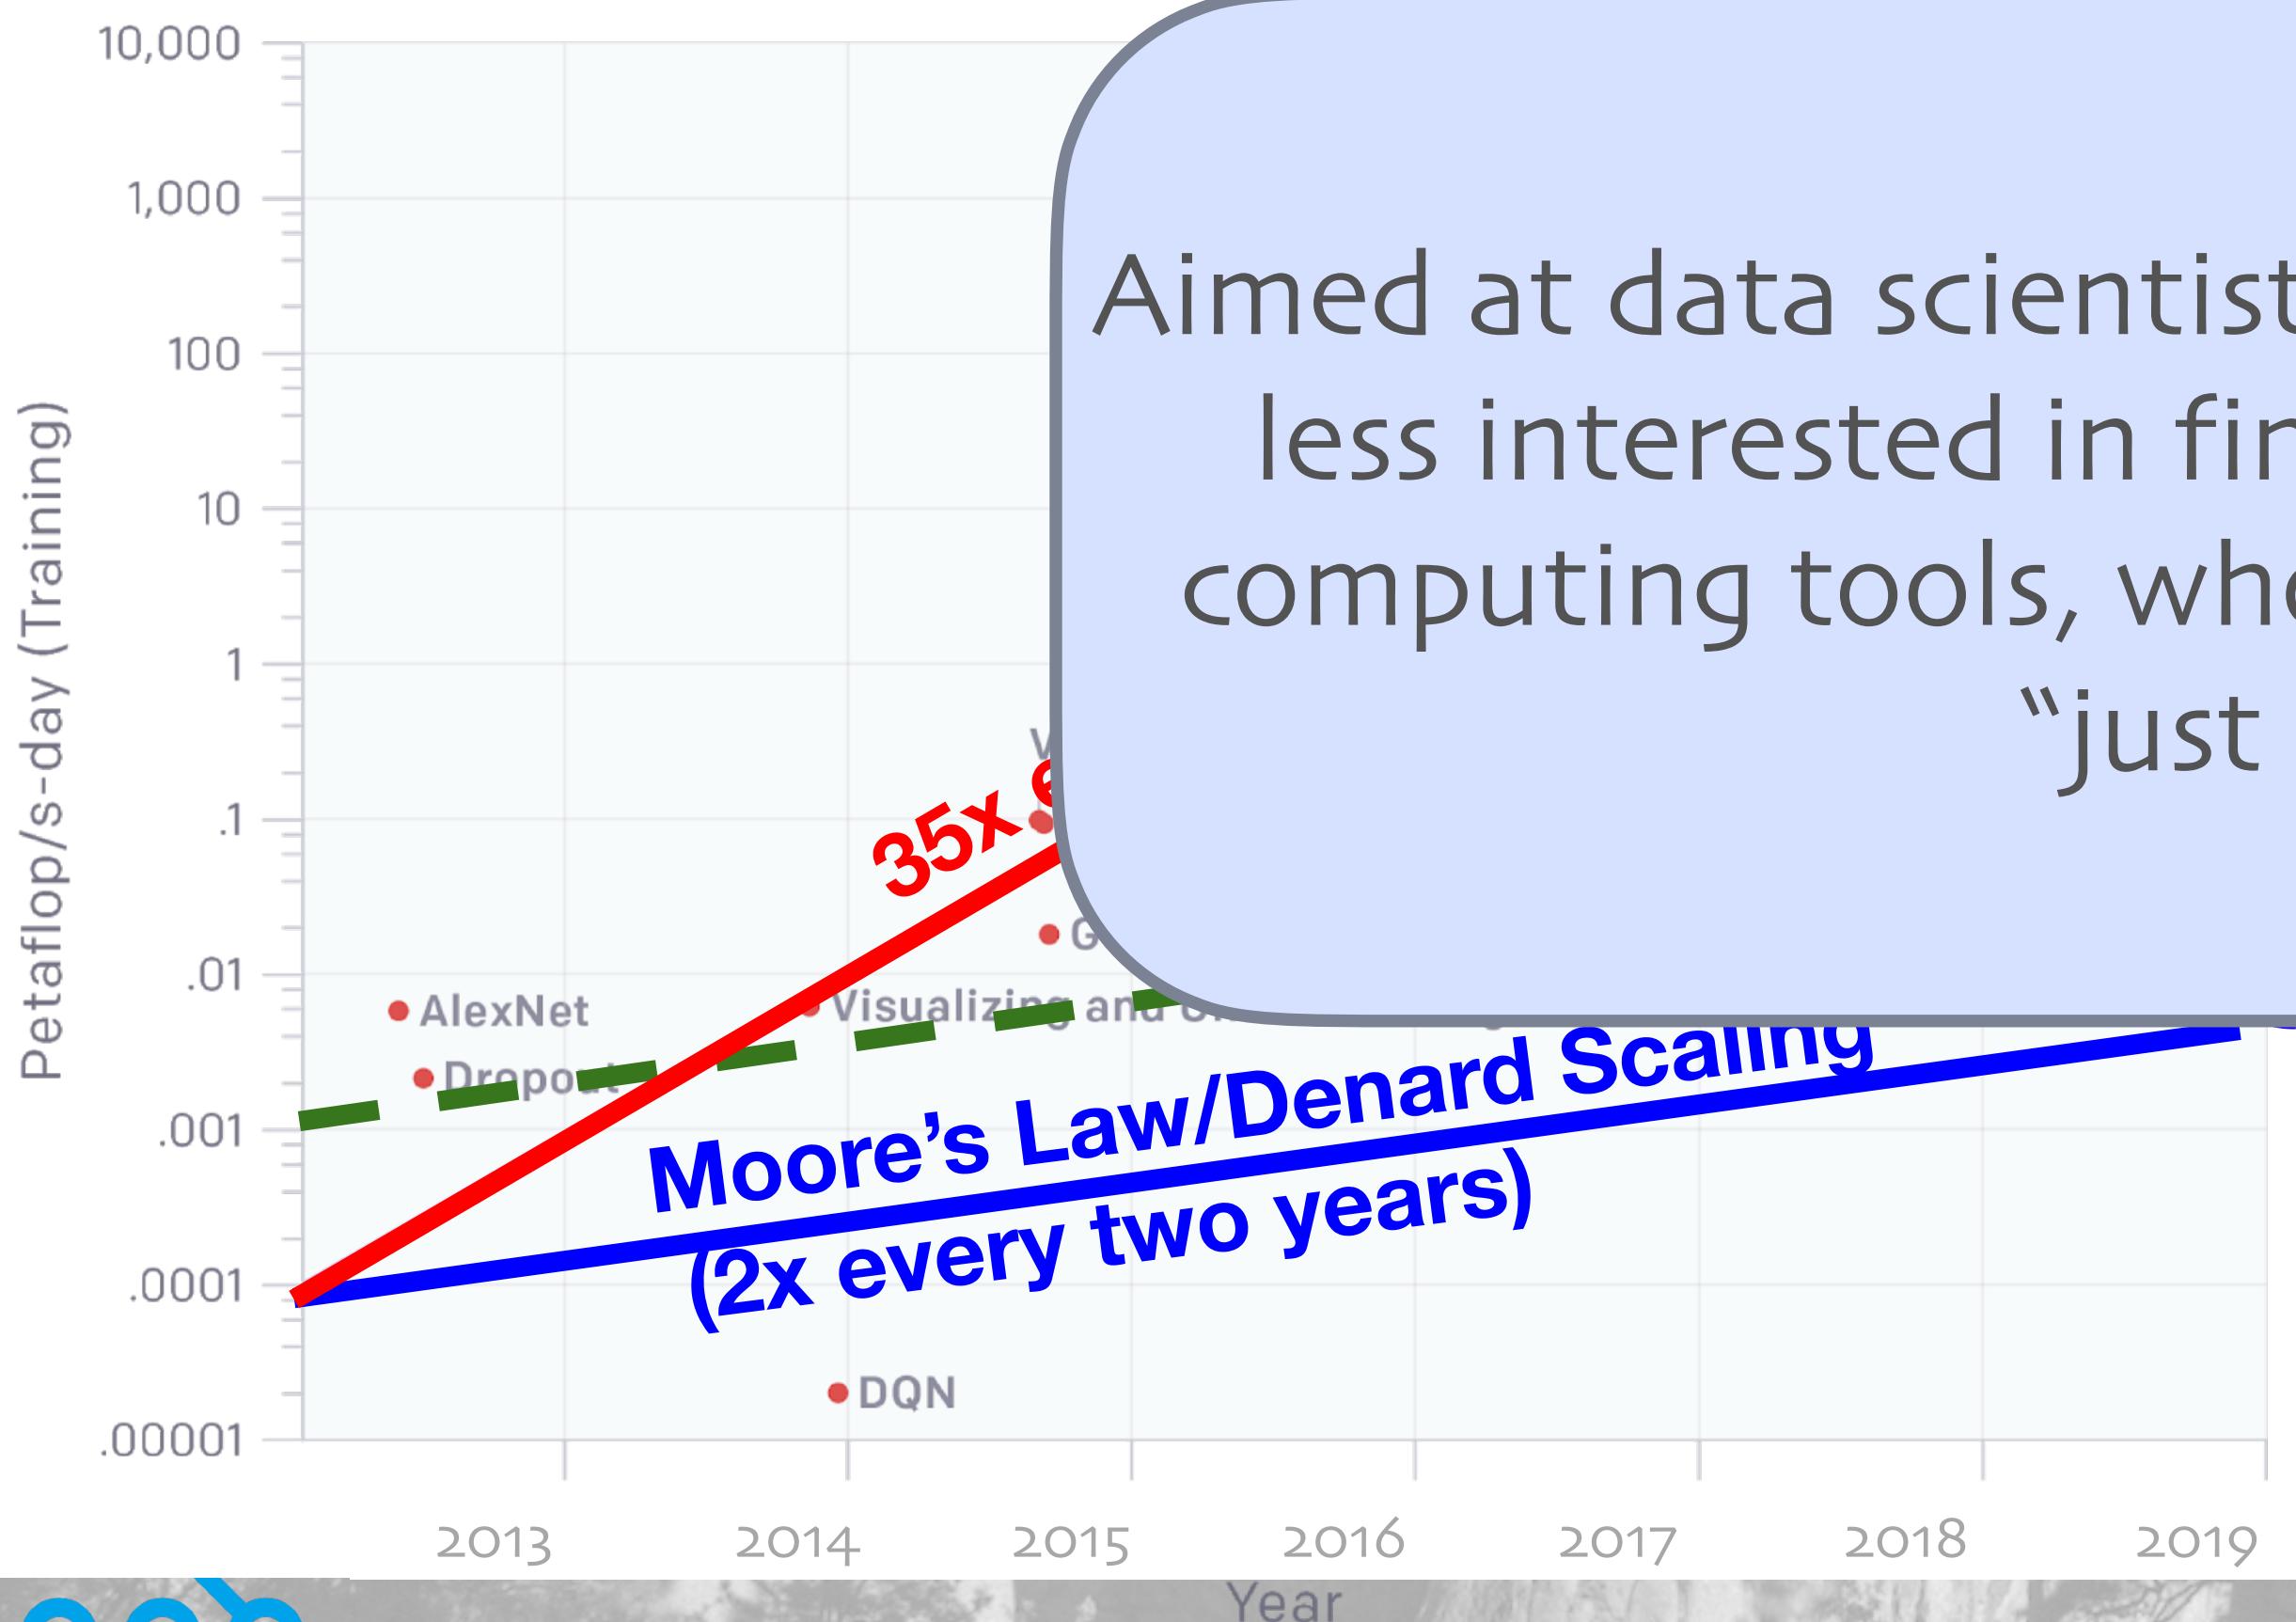
Two Major Trends

Model sizes and therefore compute requirements outstripping Moore's Law

Hence, there is a pressing need for robust, easy to use solutions for distributed Python

Python growth driven by ML/AI and other data science workloads

Aimed at data scientists and engineers who are less interested in fine-grained distributed computing tools, who want something that "just works".



The ML Landscape Today

All require distributed implementations to scale

Featurization



Streaming



Hyperparam
Tuning



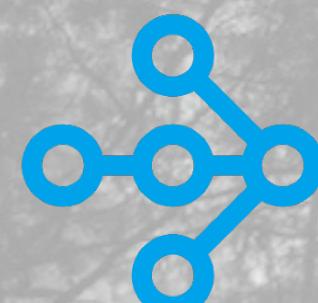
Training



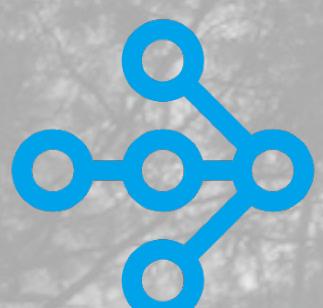
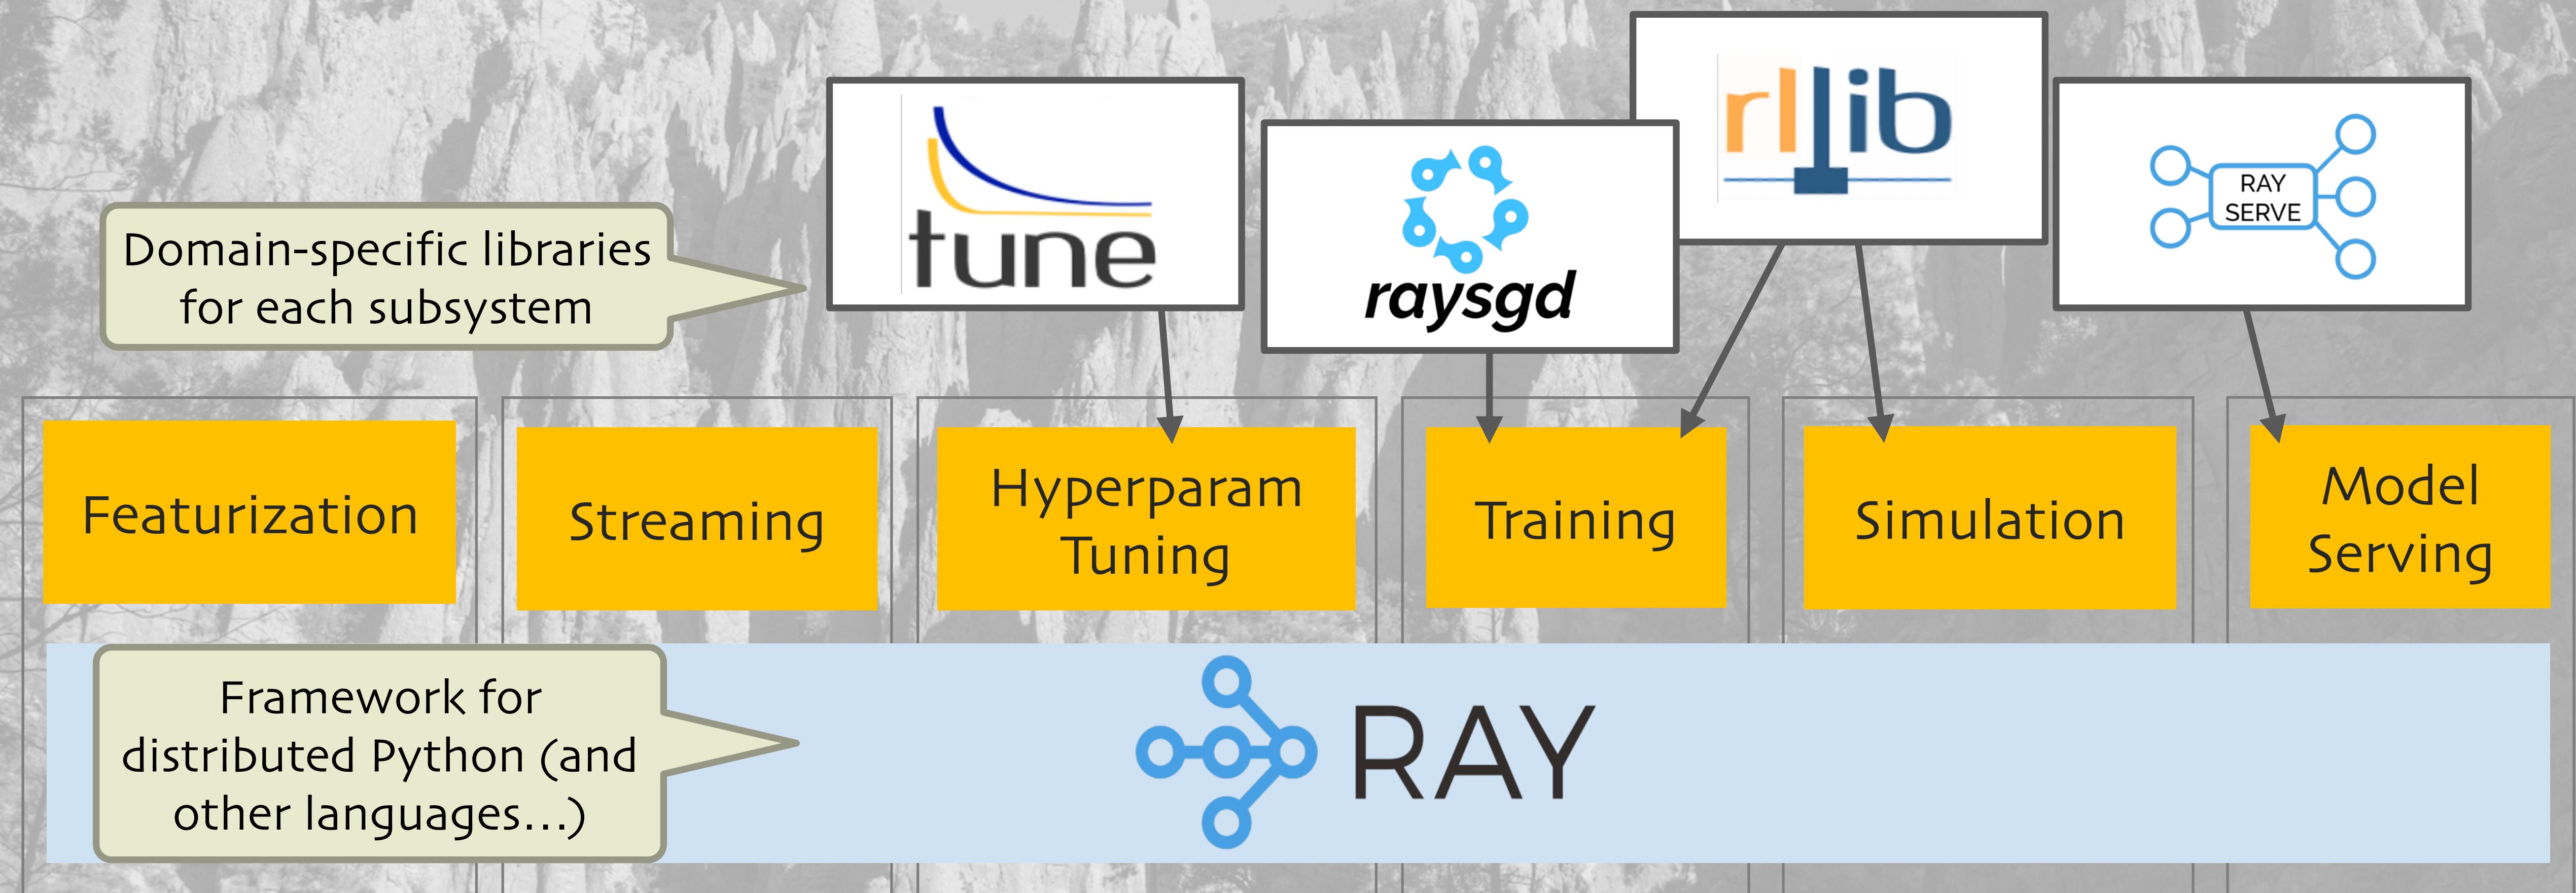
Simulation



Model
Serving



The Ray Vision: Sharing a Common Framework



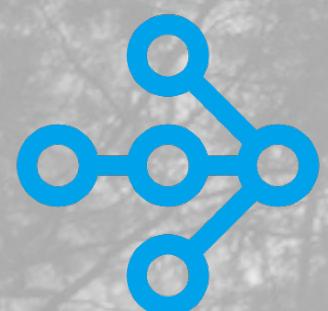
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
def add_arrays(a, b):  
    return np.add(a, b)
```

The Python you
already know...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

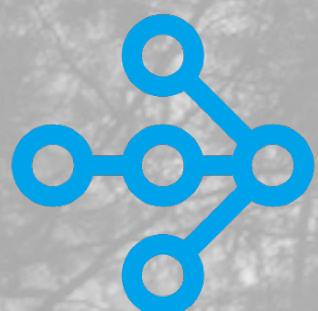
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

For completeness, add these first:

```
import ray  
import numpy as np  
ray.init()
```

Now these functions
are remote "tasks"



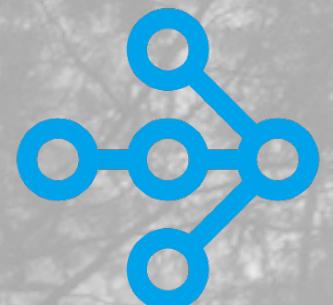
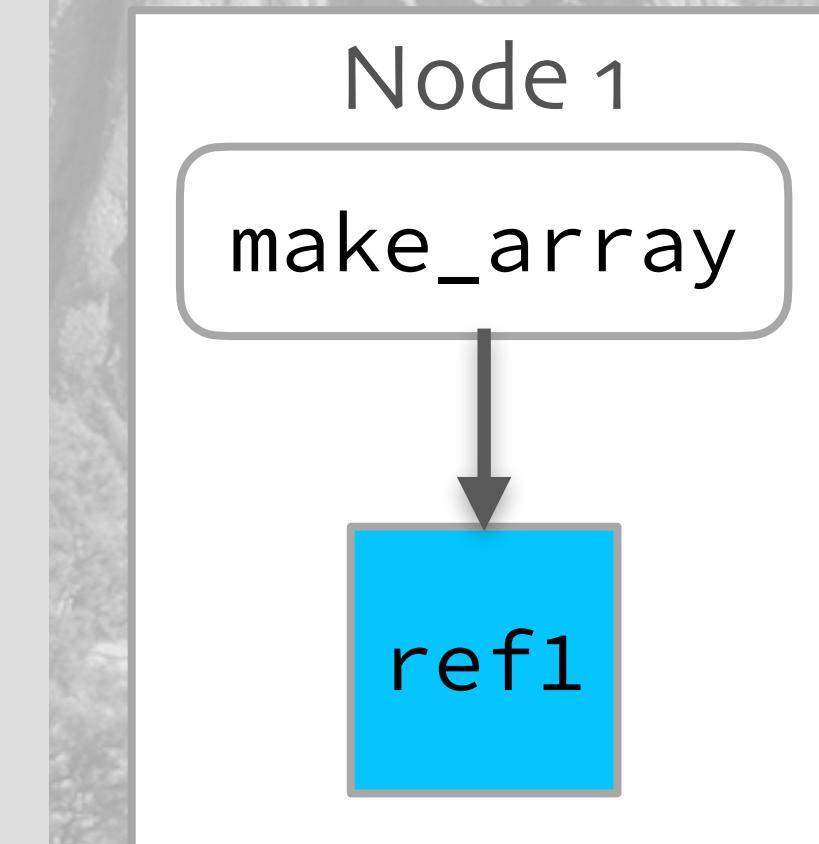
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
ref1 = make_array.remote(...)
```



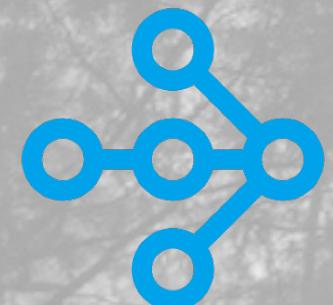
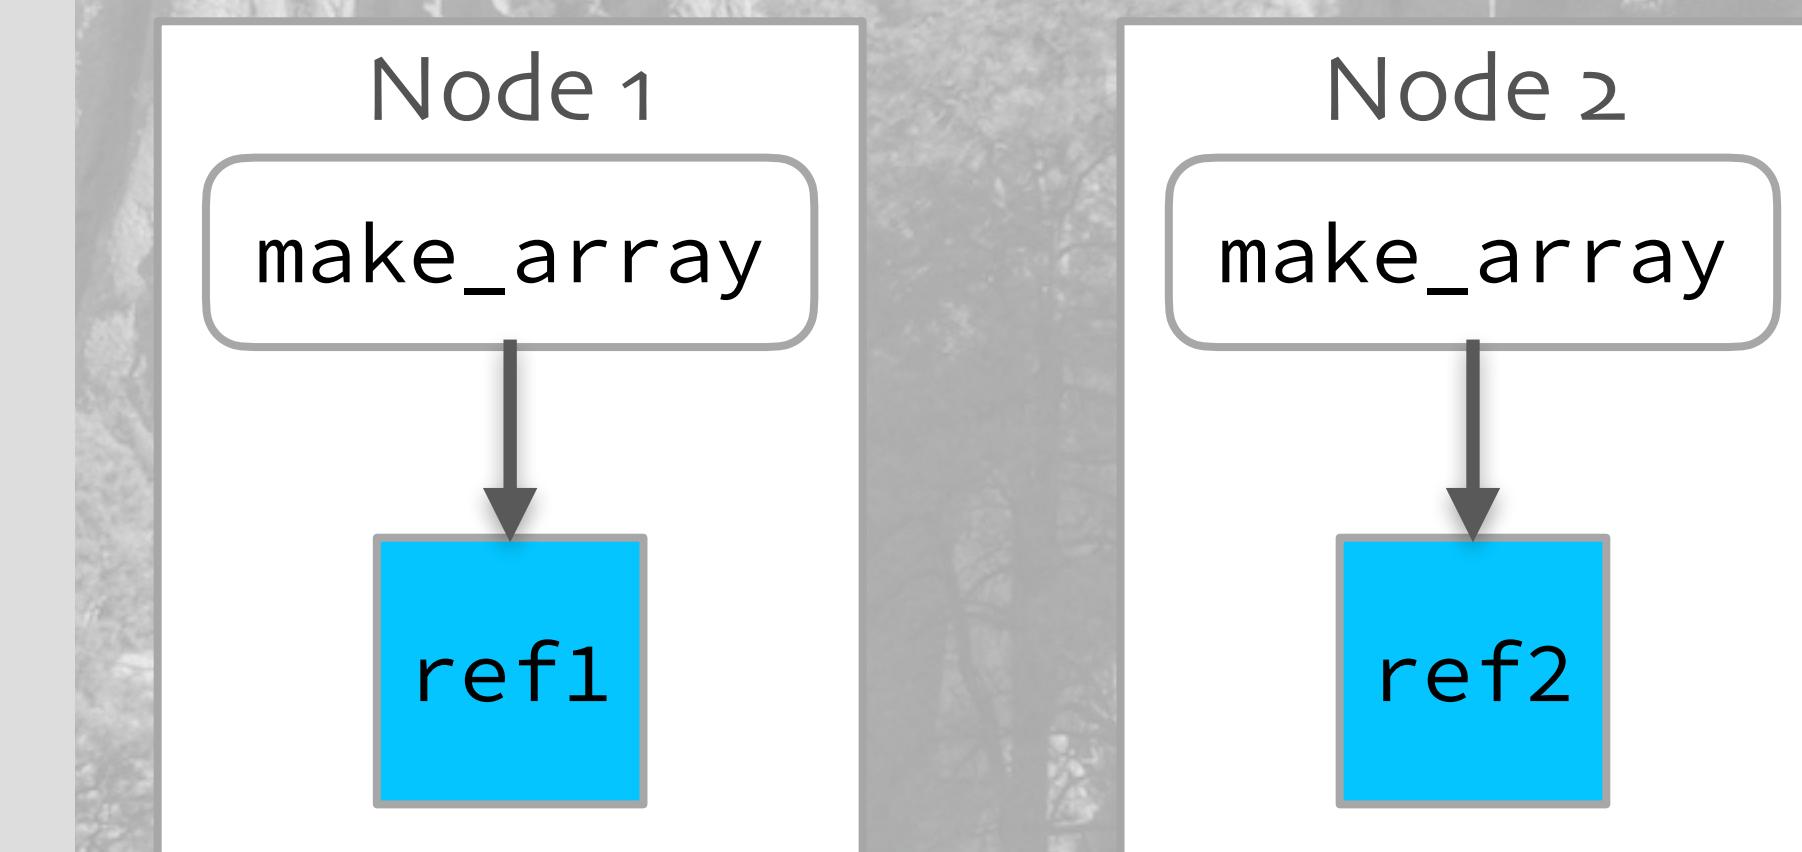
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

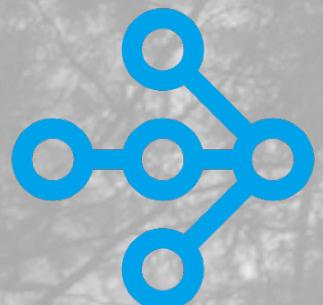
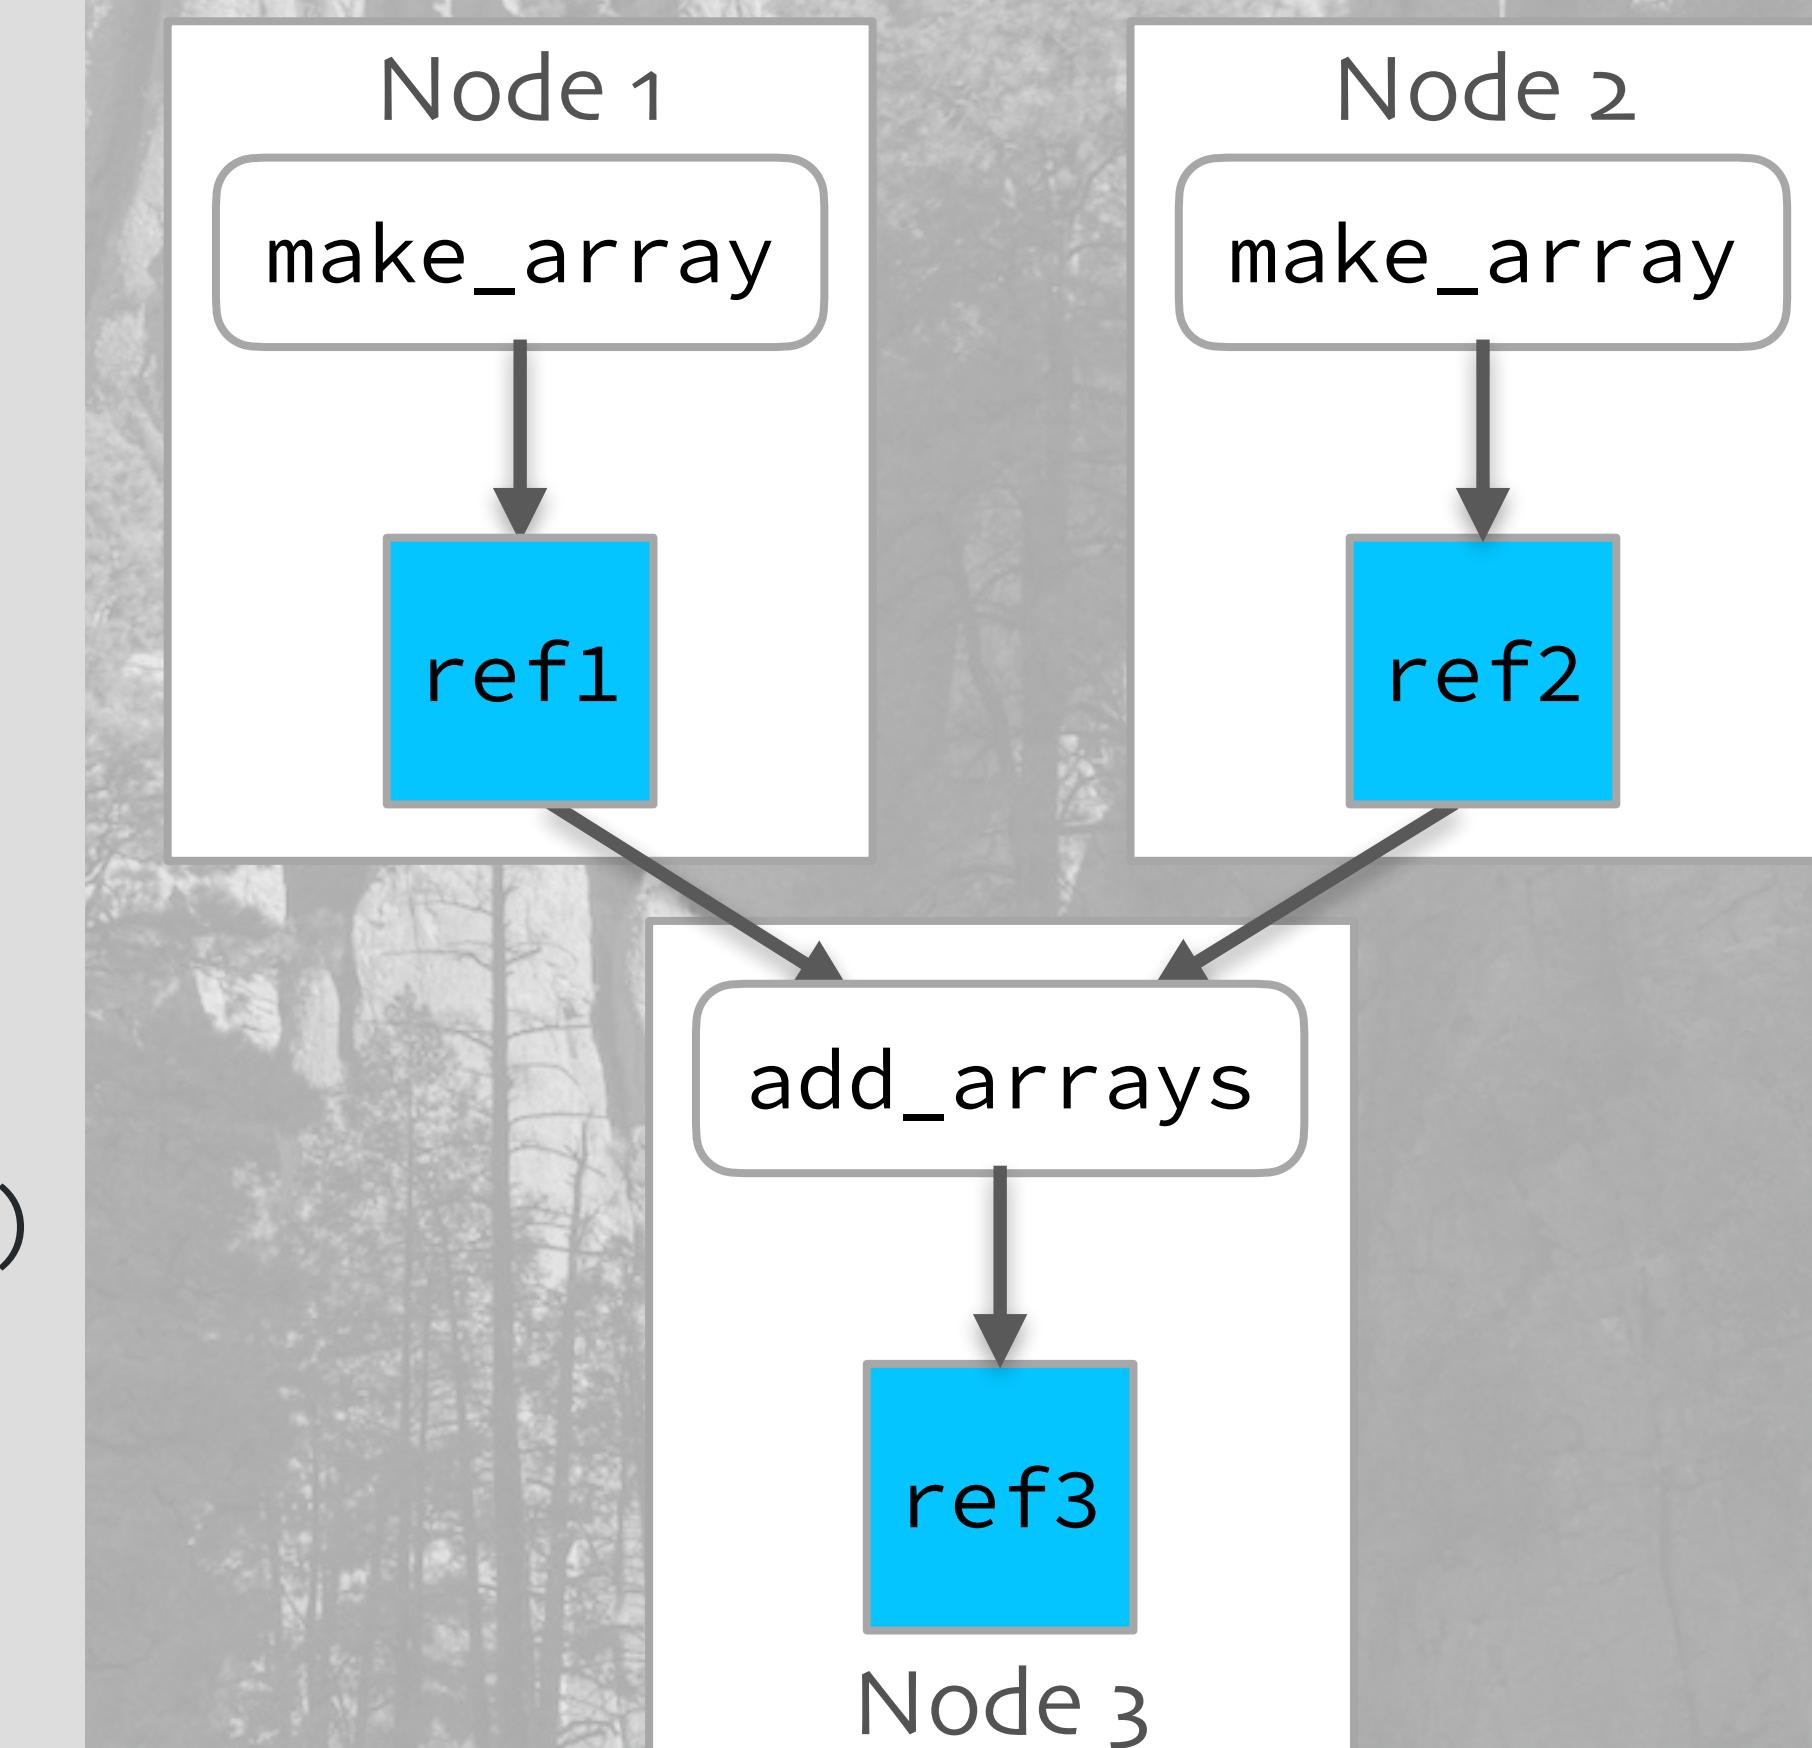
```
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

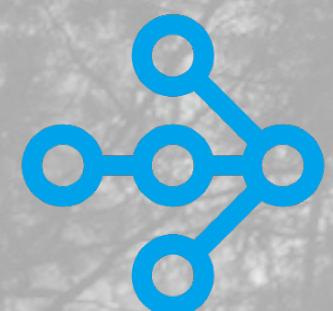
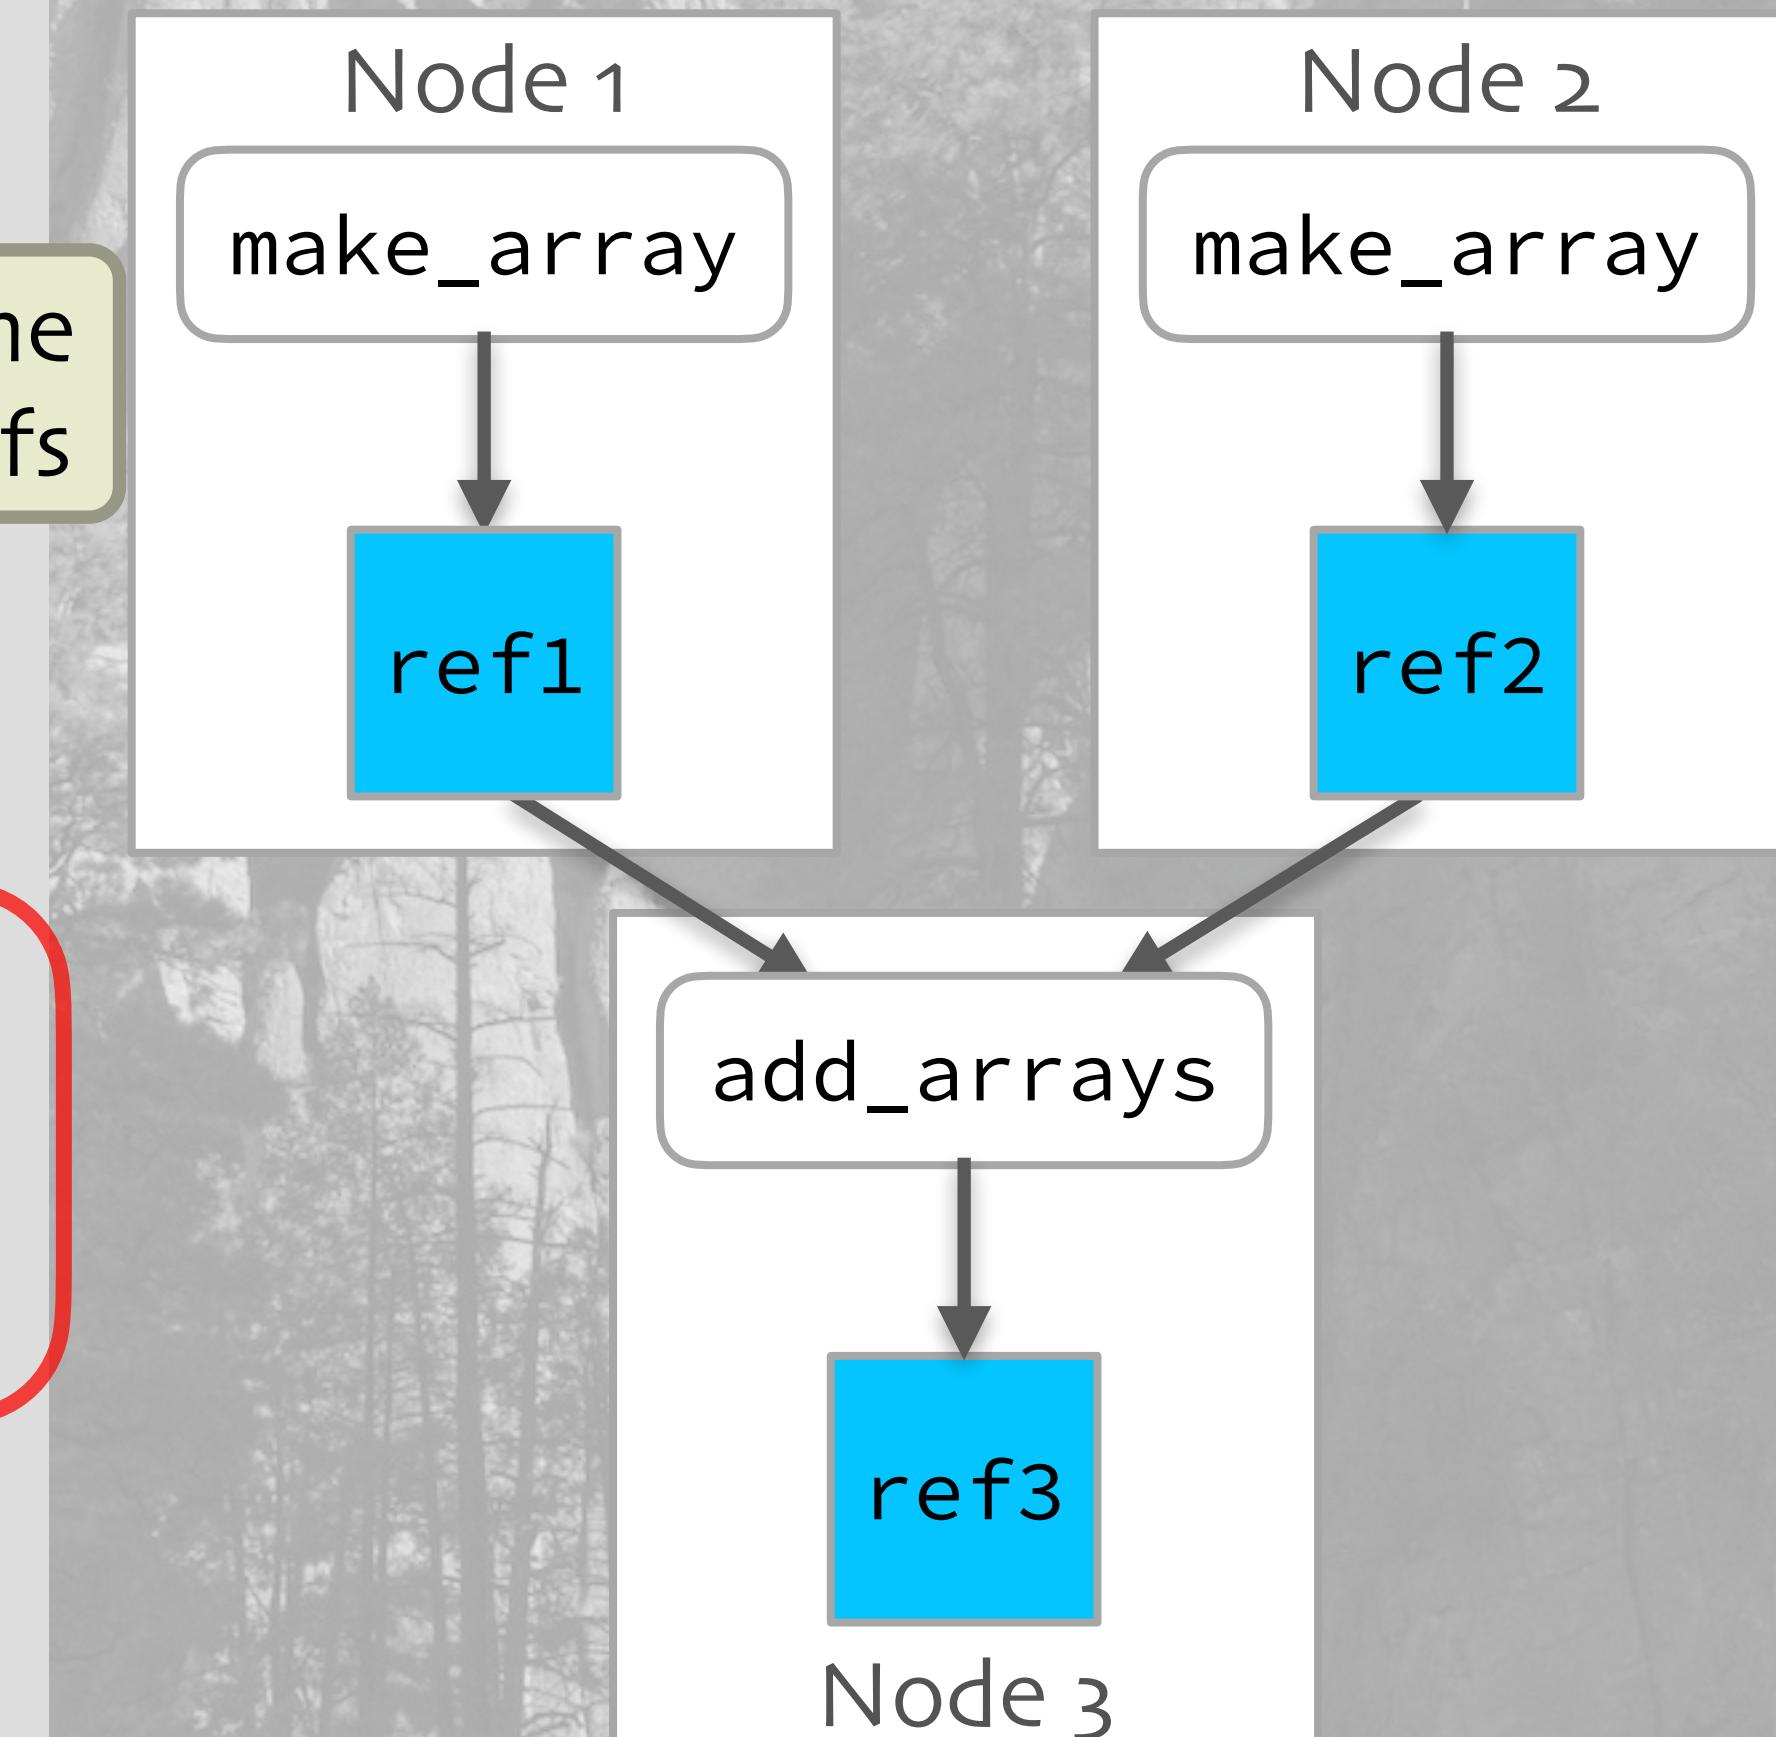
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

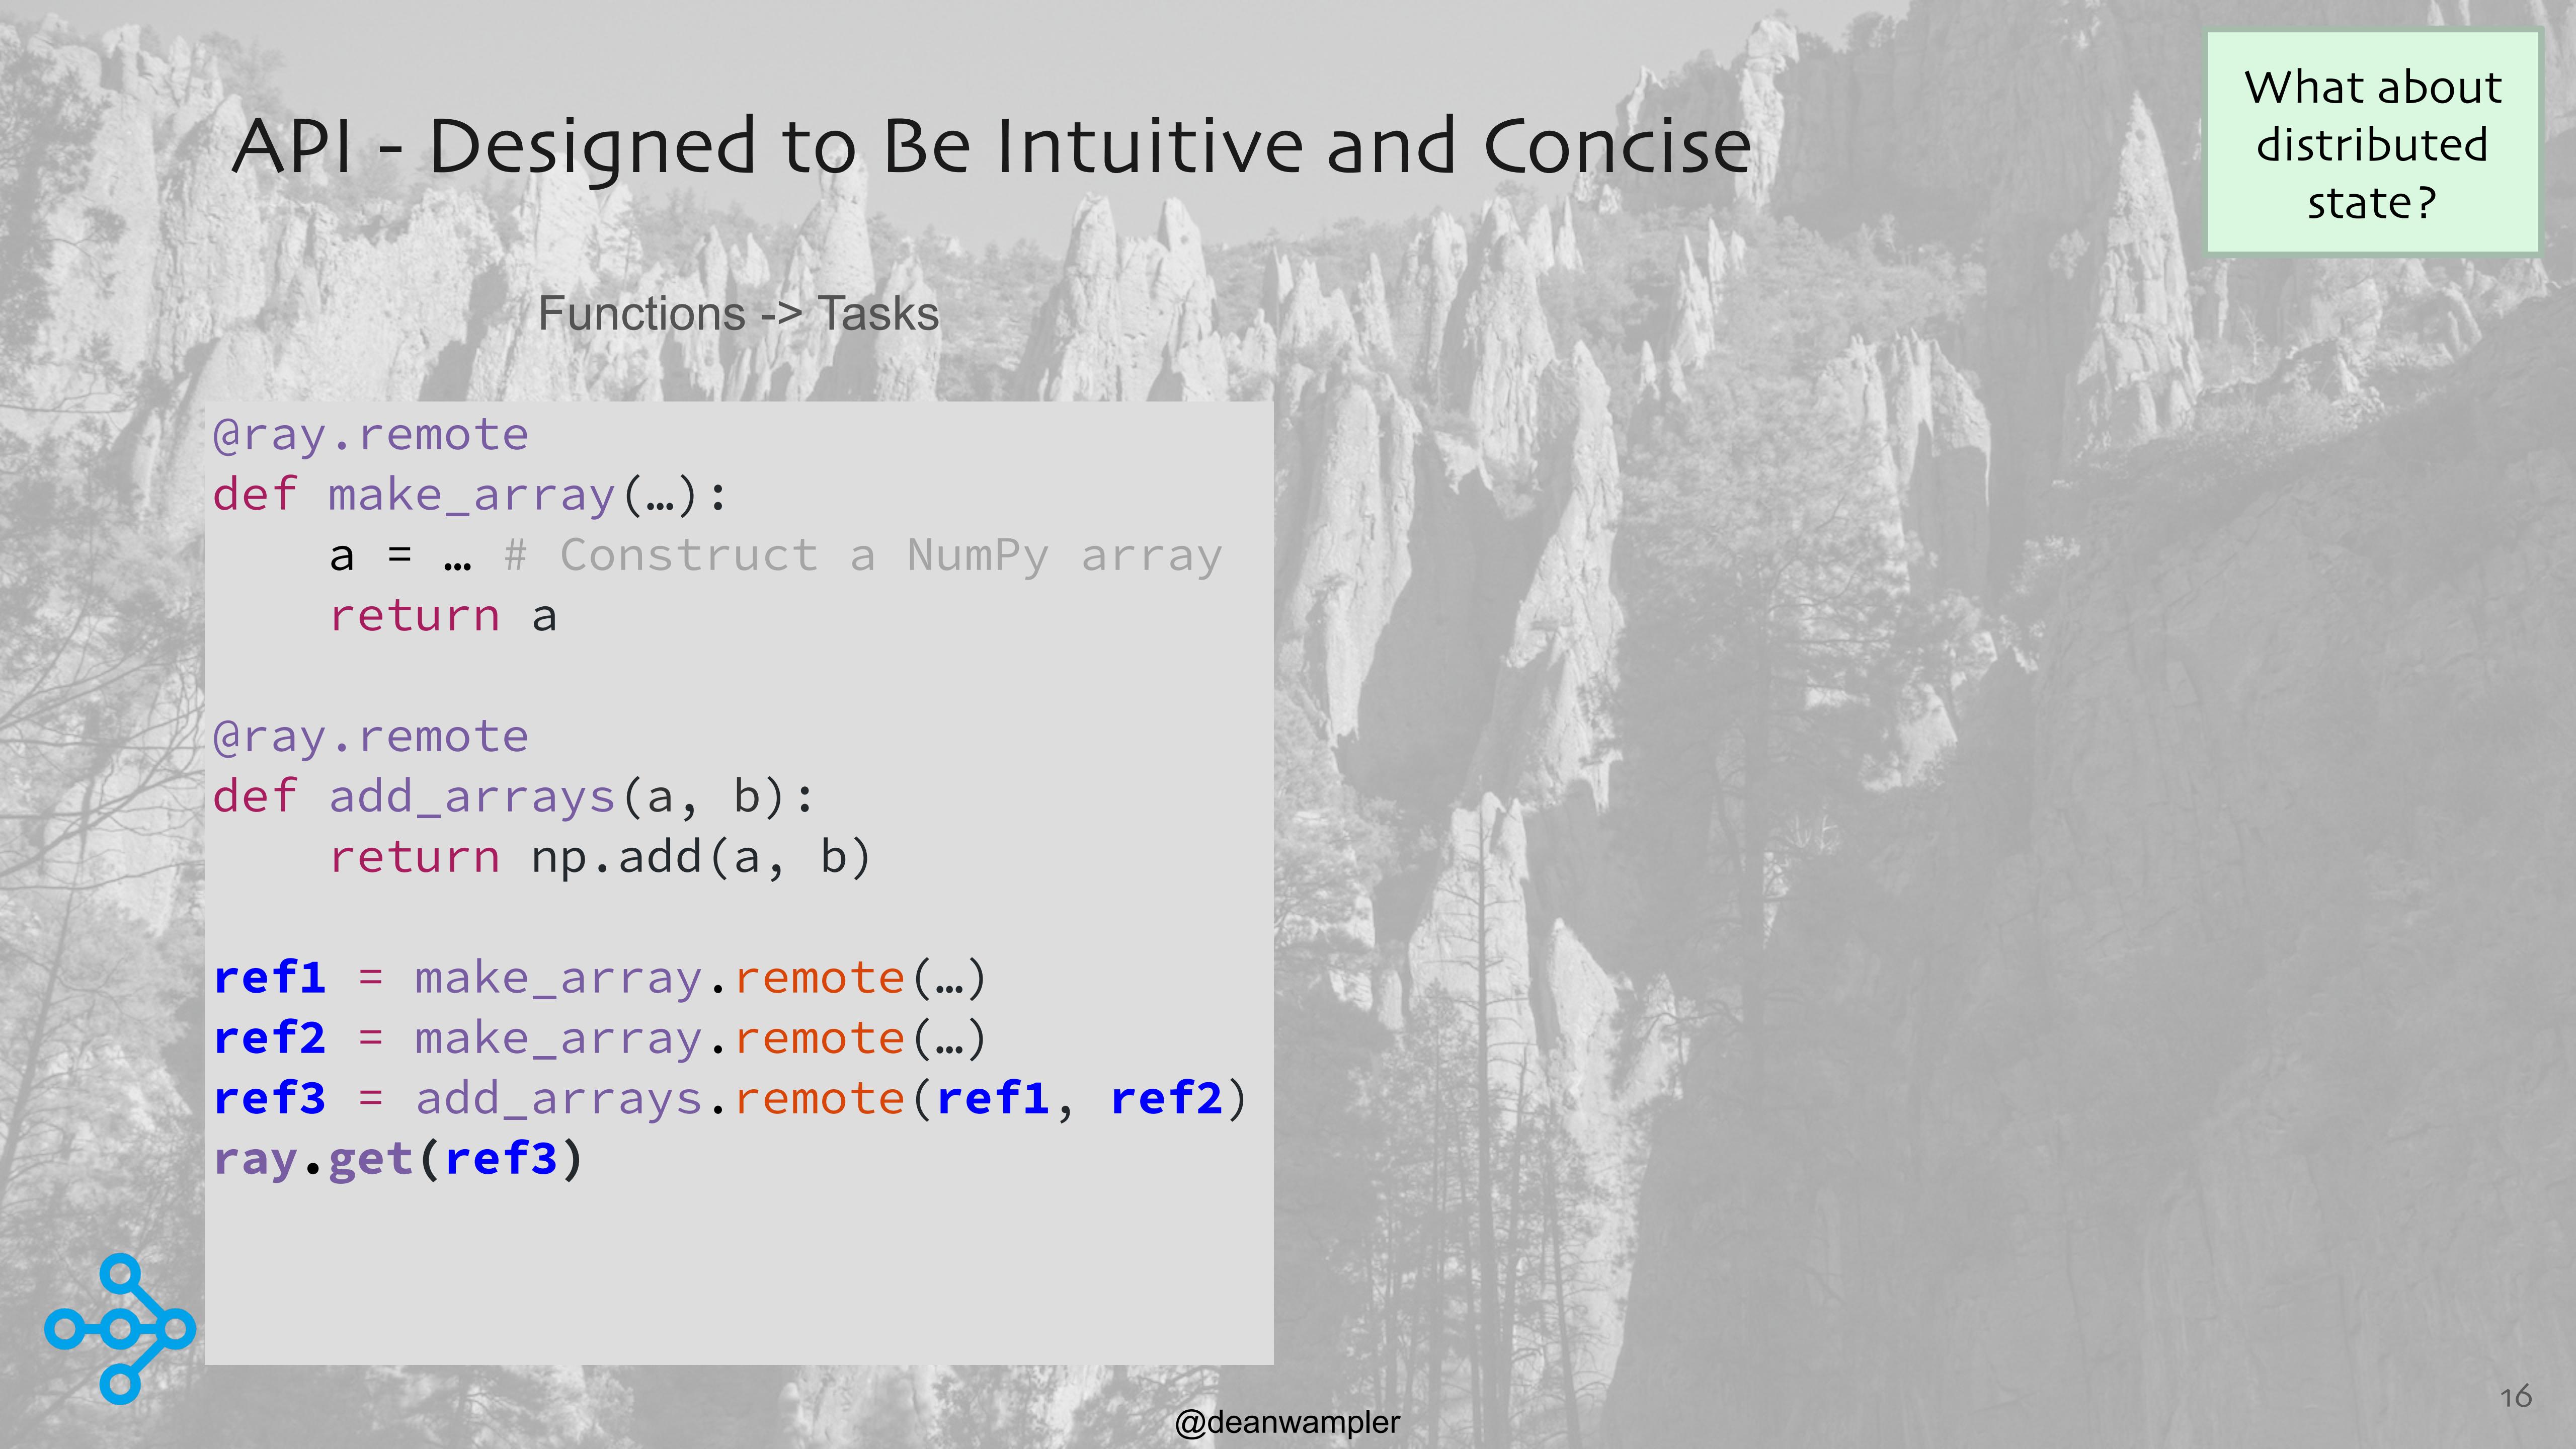
```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```

Ray handles sequencing
of async dependencies

Ray handles extracting the
arrays from the object refs





What about
distributed
state?

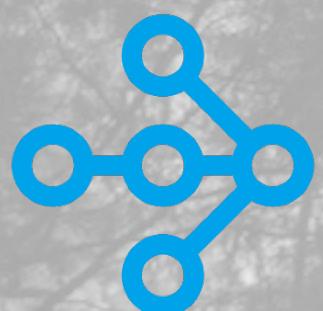
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

ref1 = make_array.remote(...)
ref2 = make_array.remote(...)
ref3 = add_arrays.remote(ref1, ref2)
ray.get(ref3)
```



API - Designed to Be Intuitive and Concise

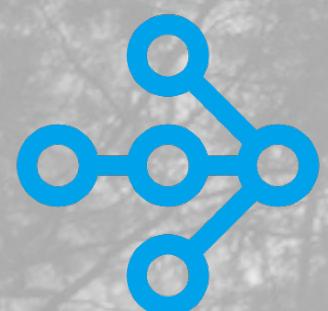
Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```

Classes -> Actors

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
    return self.value
```

The Python
classes you
love...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

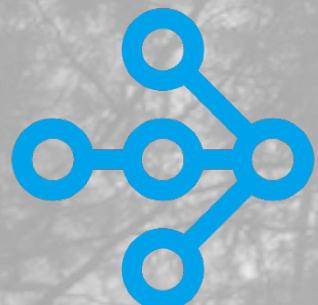
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```

Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
        return self.value  
    def get_count(self):  
        return self.value
```

... now a remote
“actor”

You need a
“getter” method
to read the state.



API - Designed to Be Intuitive and Concise

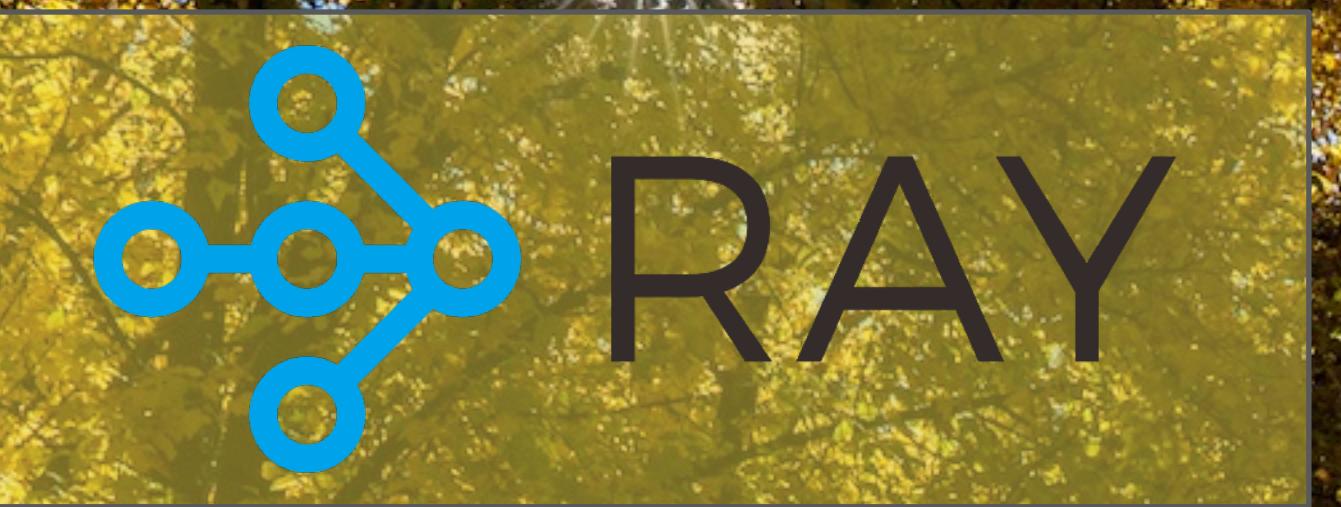
Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```



Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
        return self.value  
    def get_count(self):  
        return self.value  
  
c = Counter.remote()  
ref4 = c.increment.remote()  
ref5 = c.increment.remote()  
ray.get([ref4, ref5]) # [1, 2]
```

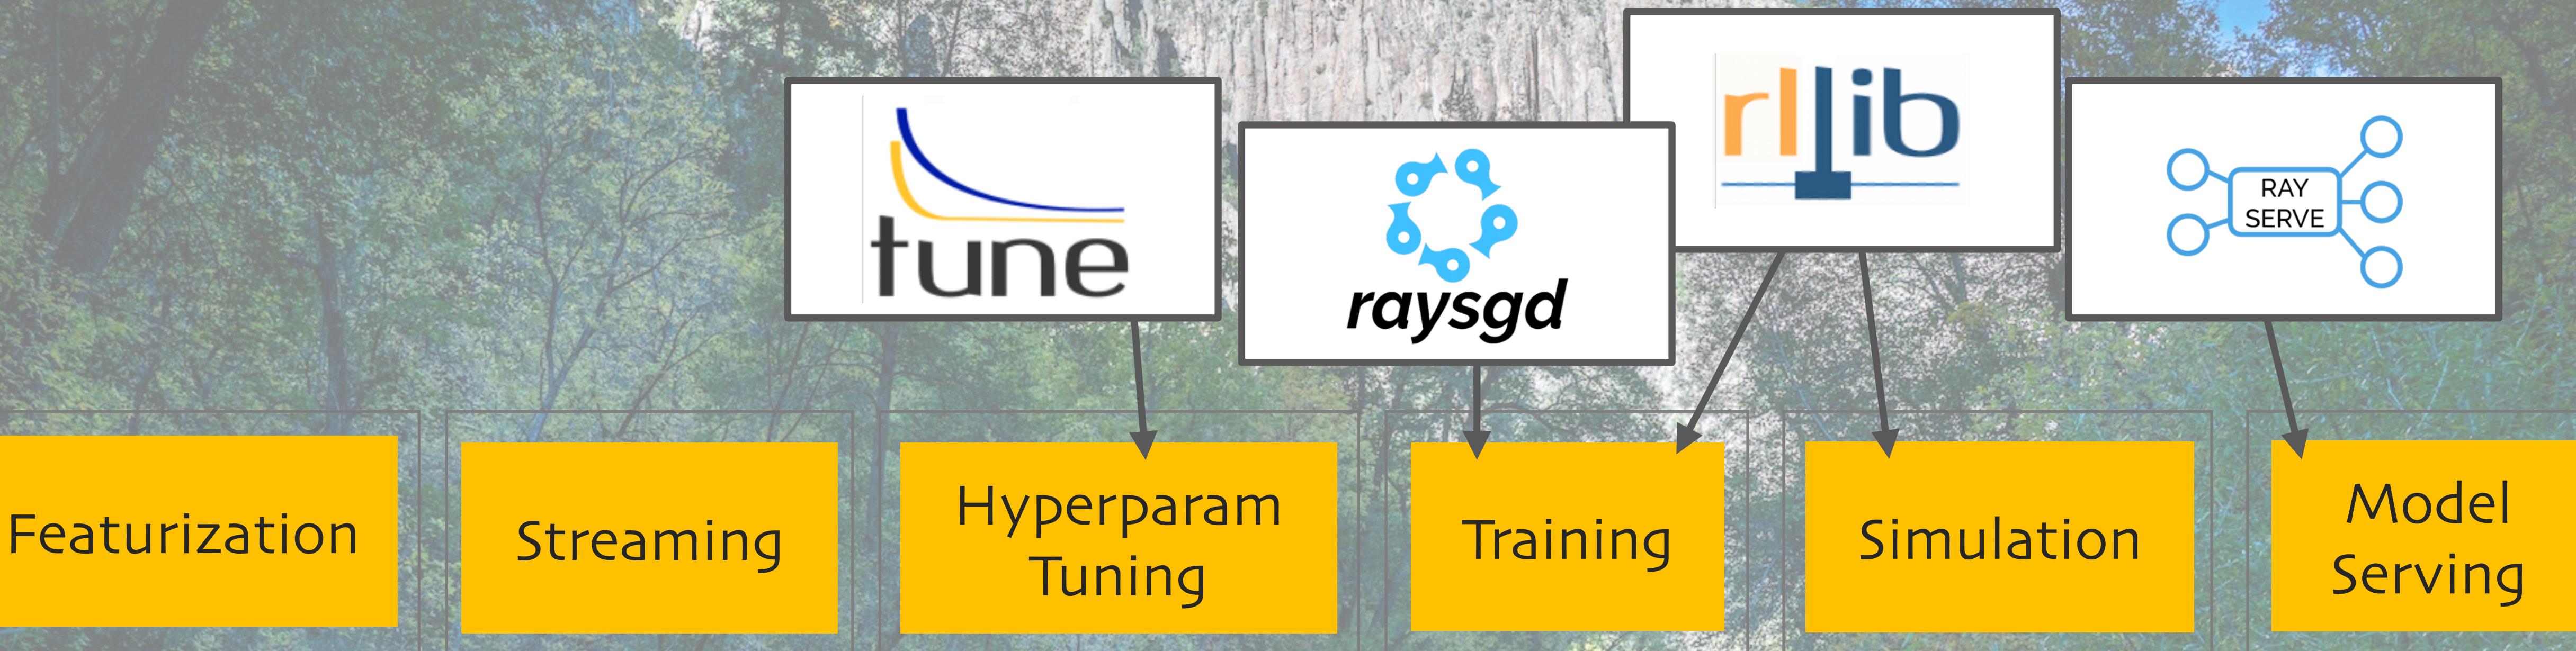


Machine Learning with Ray-based Libraries

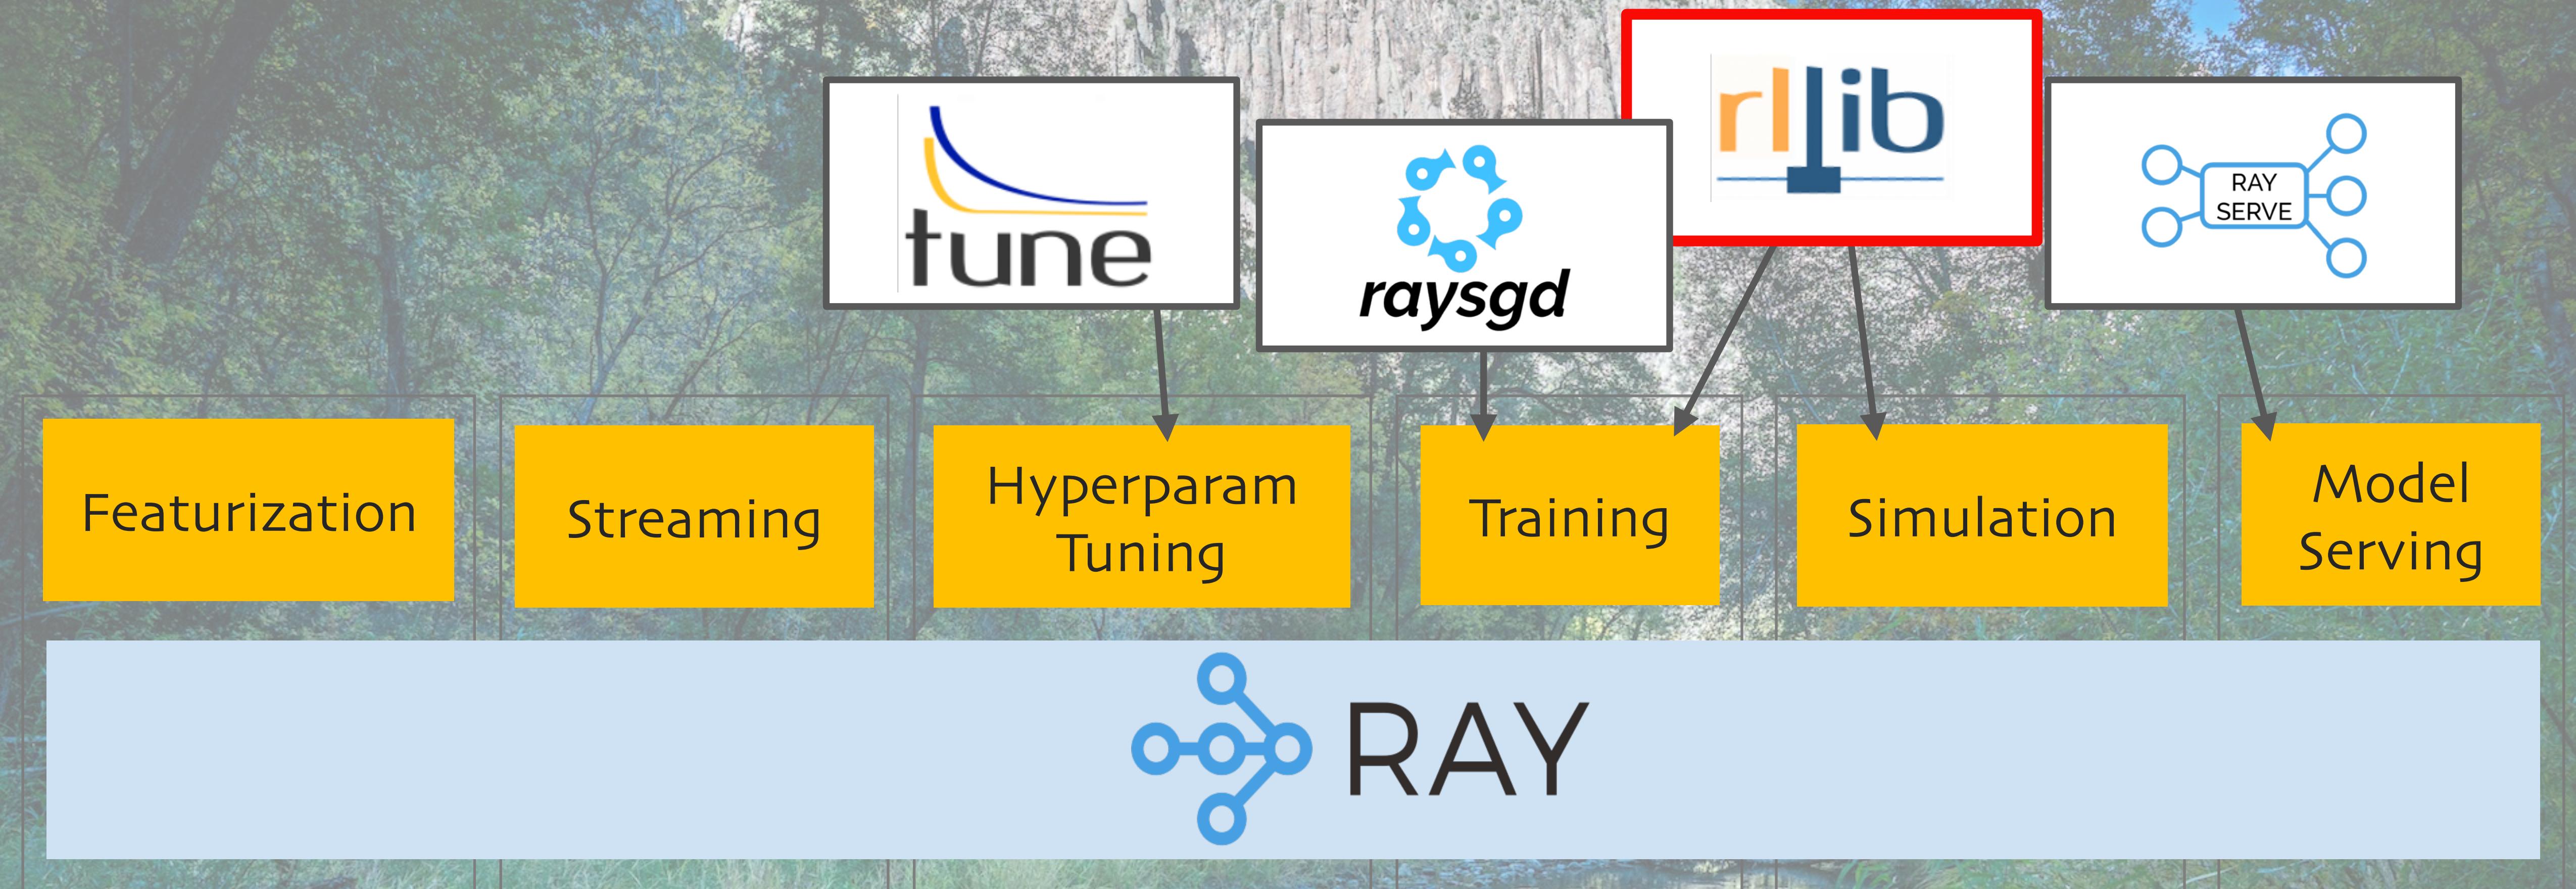


@deanwampler

Ray Libraries

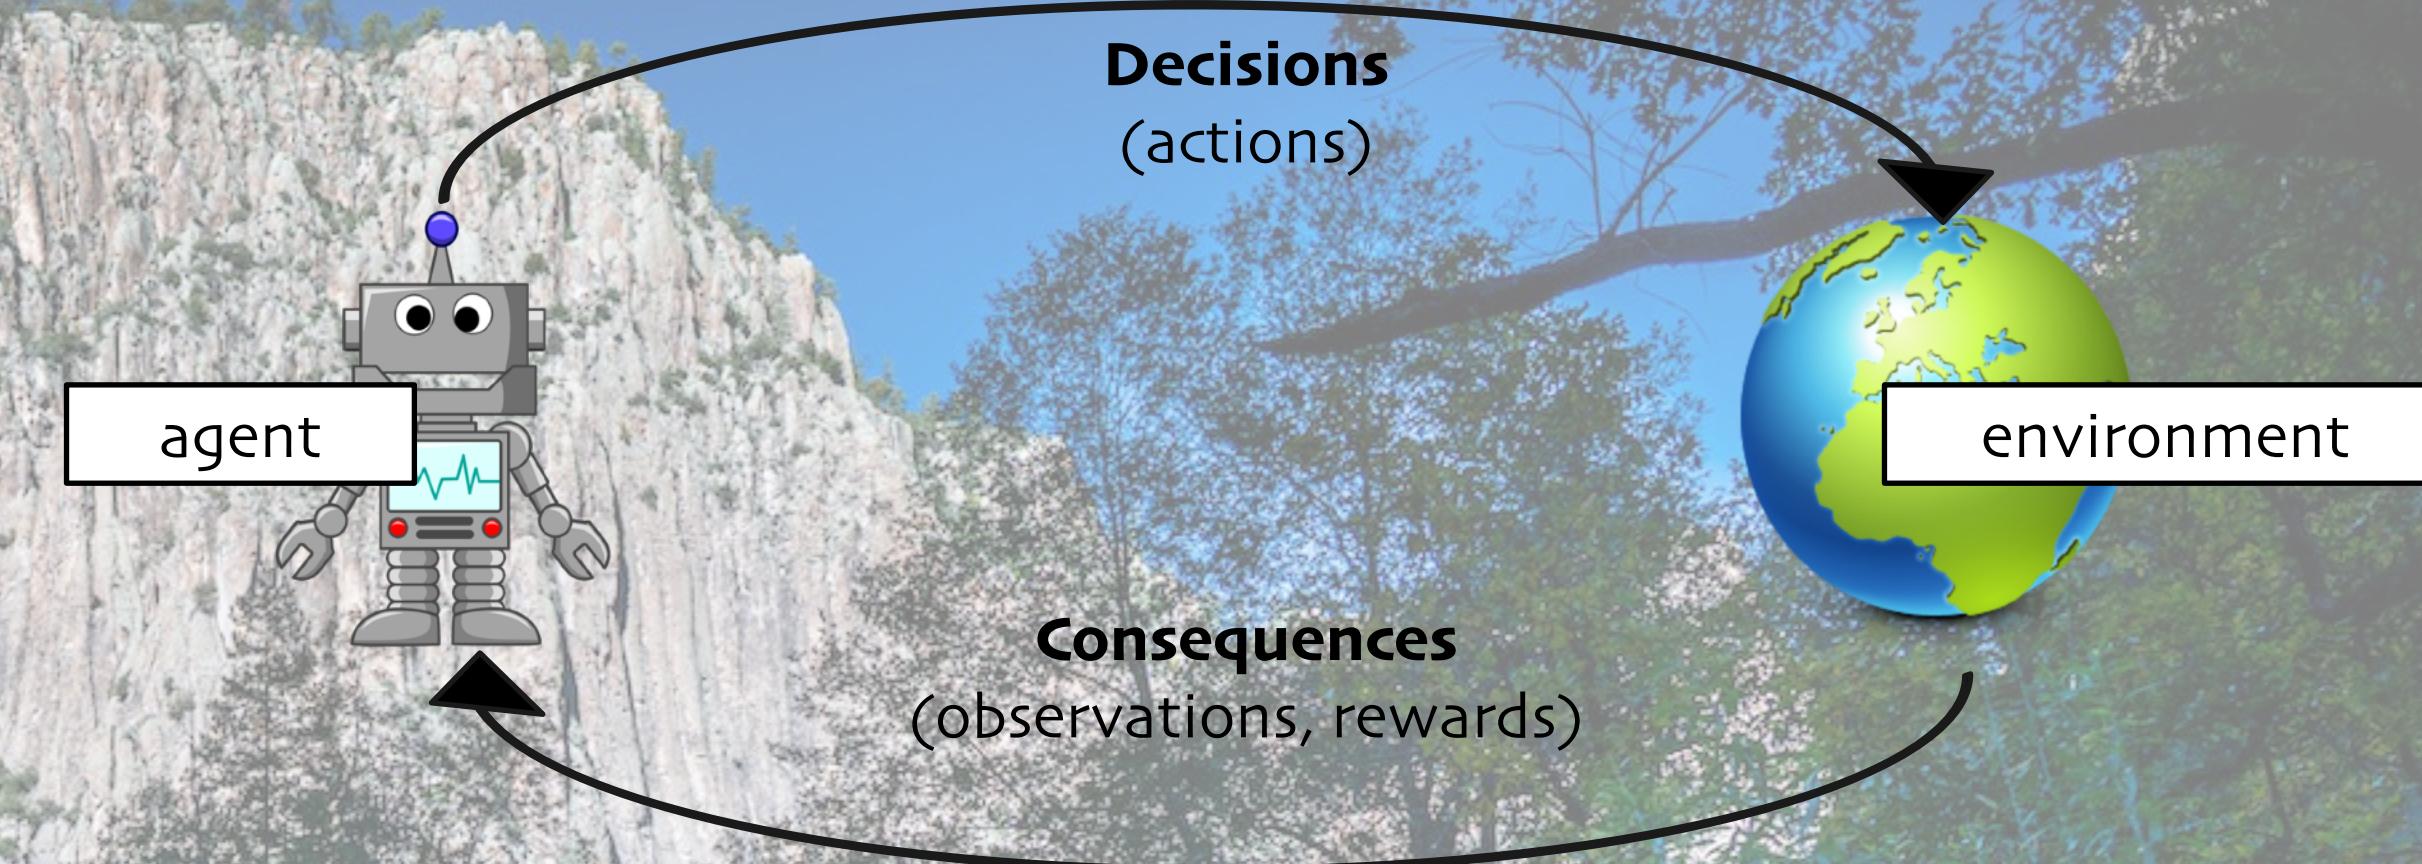


Reinforcement Learning - Ray RLlib



rllib.io

Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

Industrial
Processes

System
Optimization

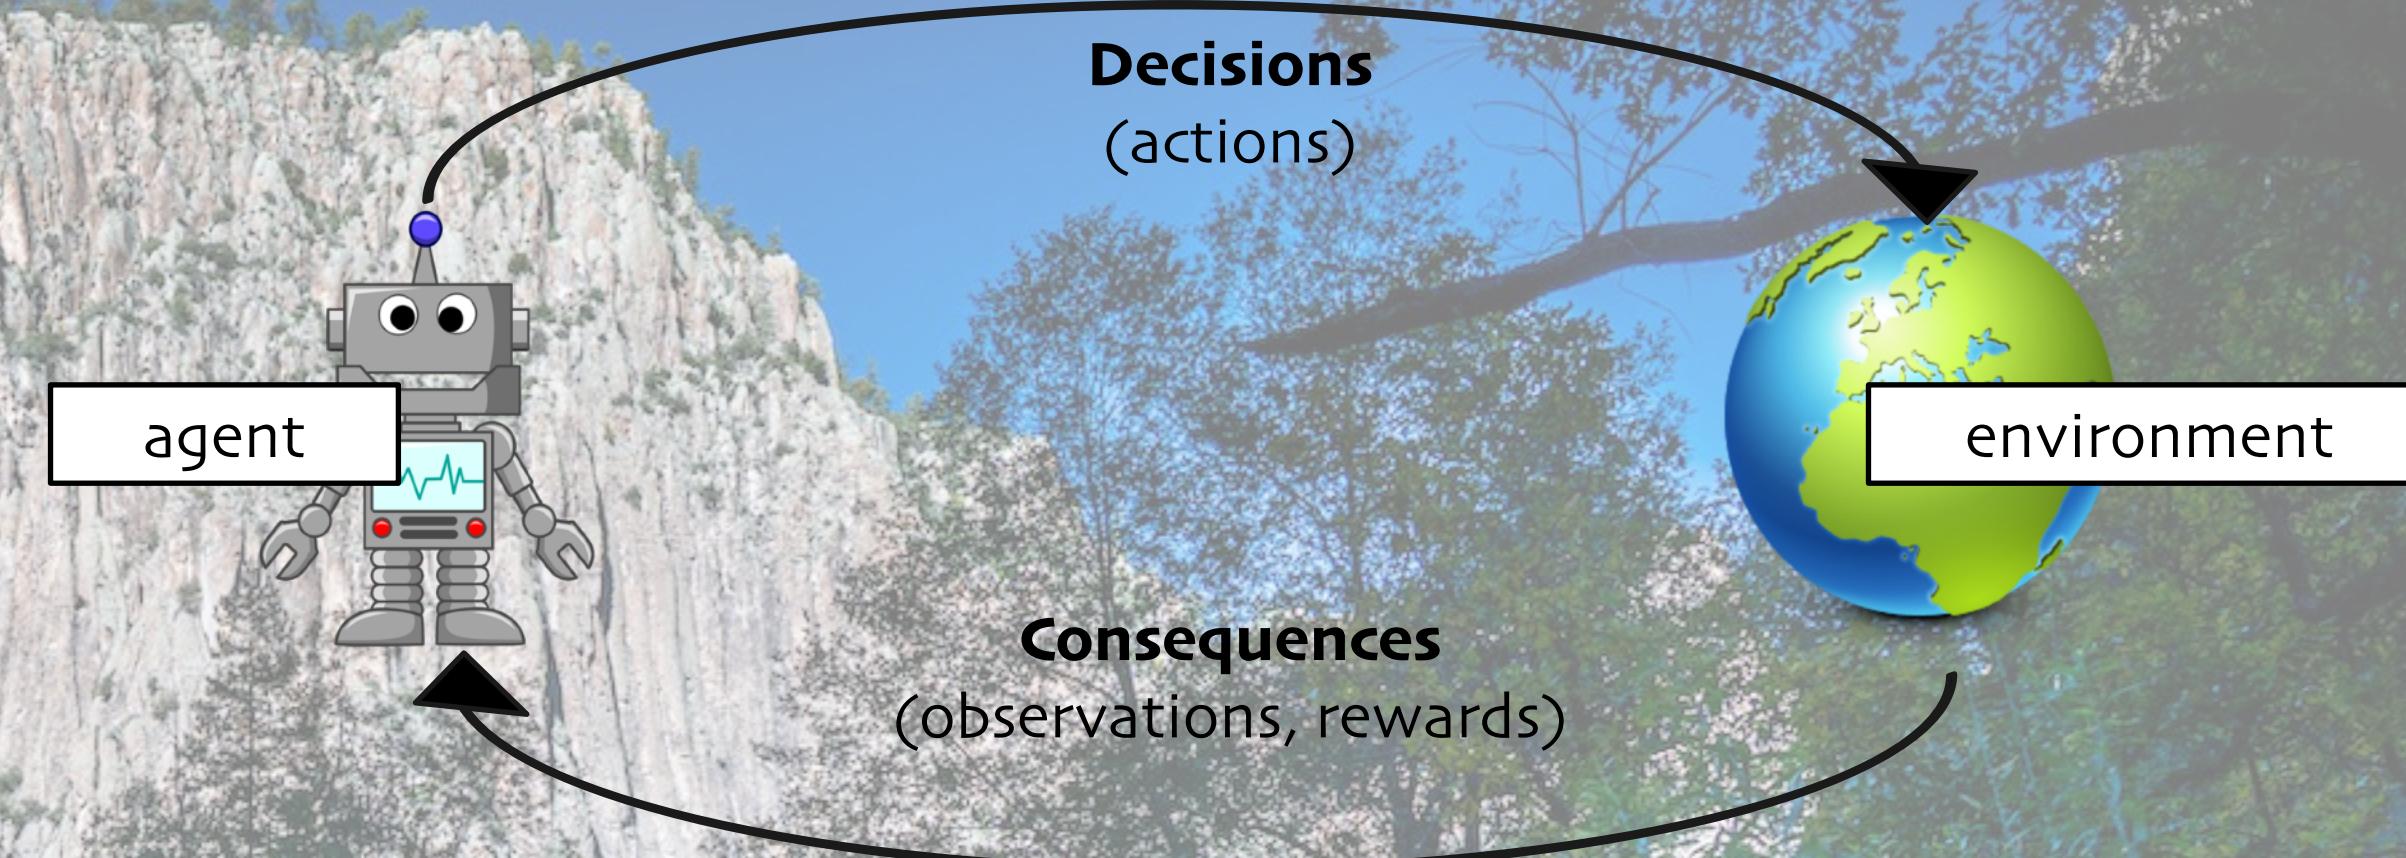
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

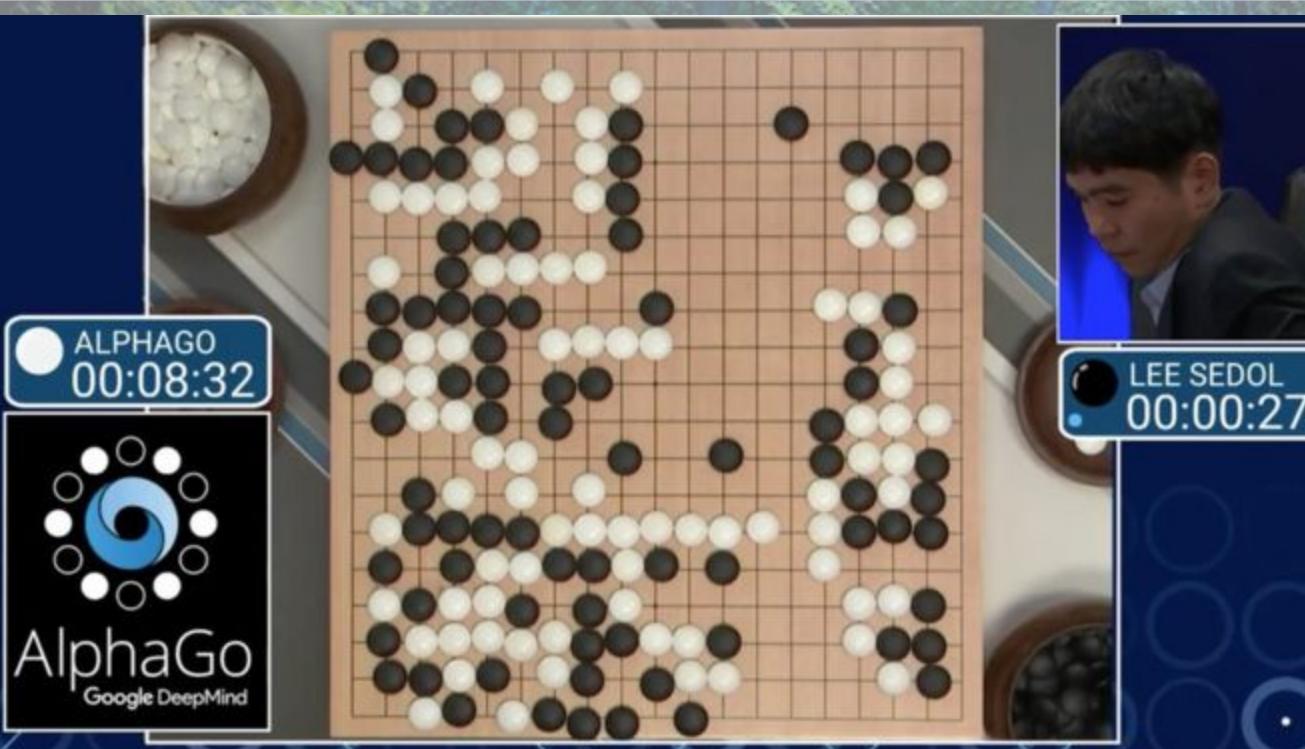
Industrial
Processes

System
Optimization

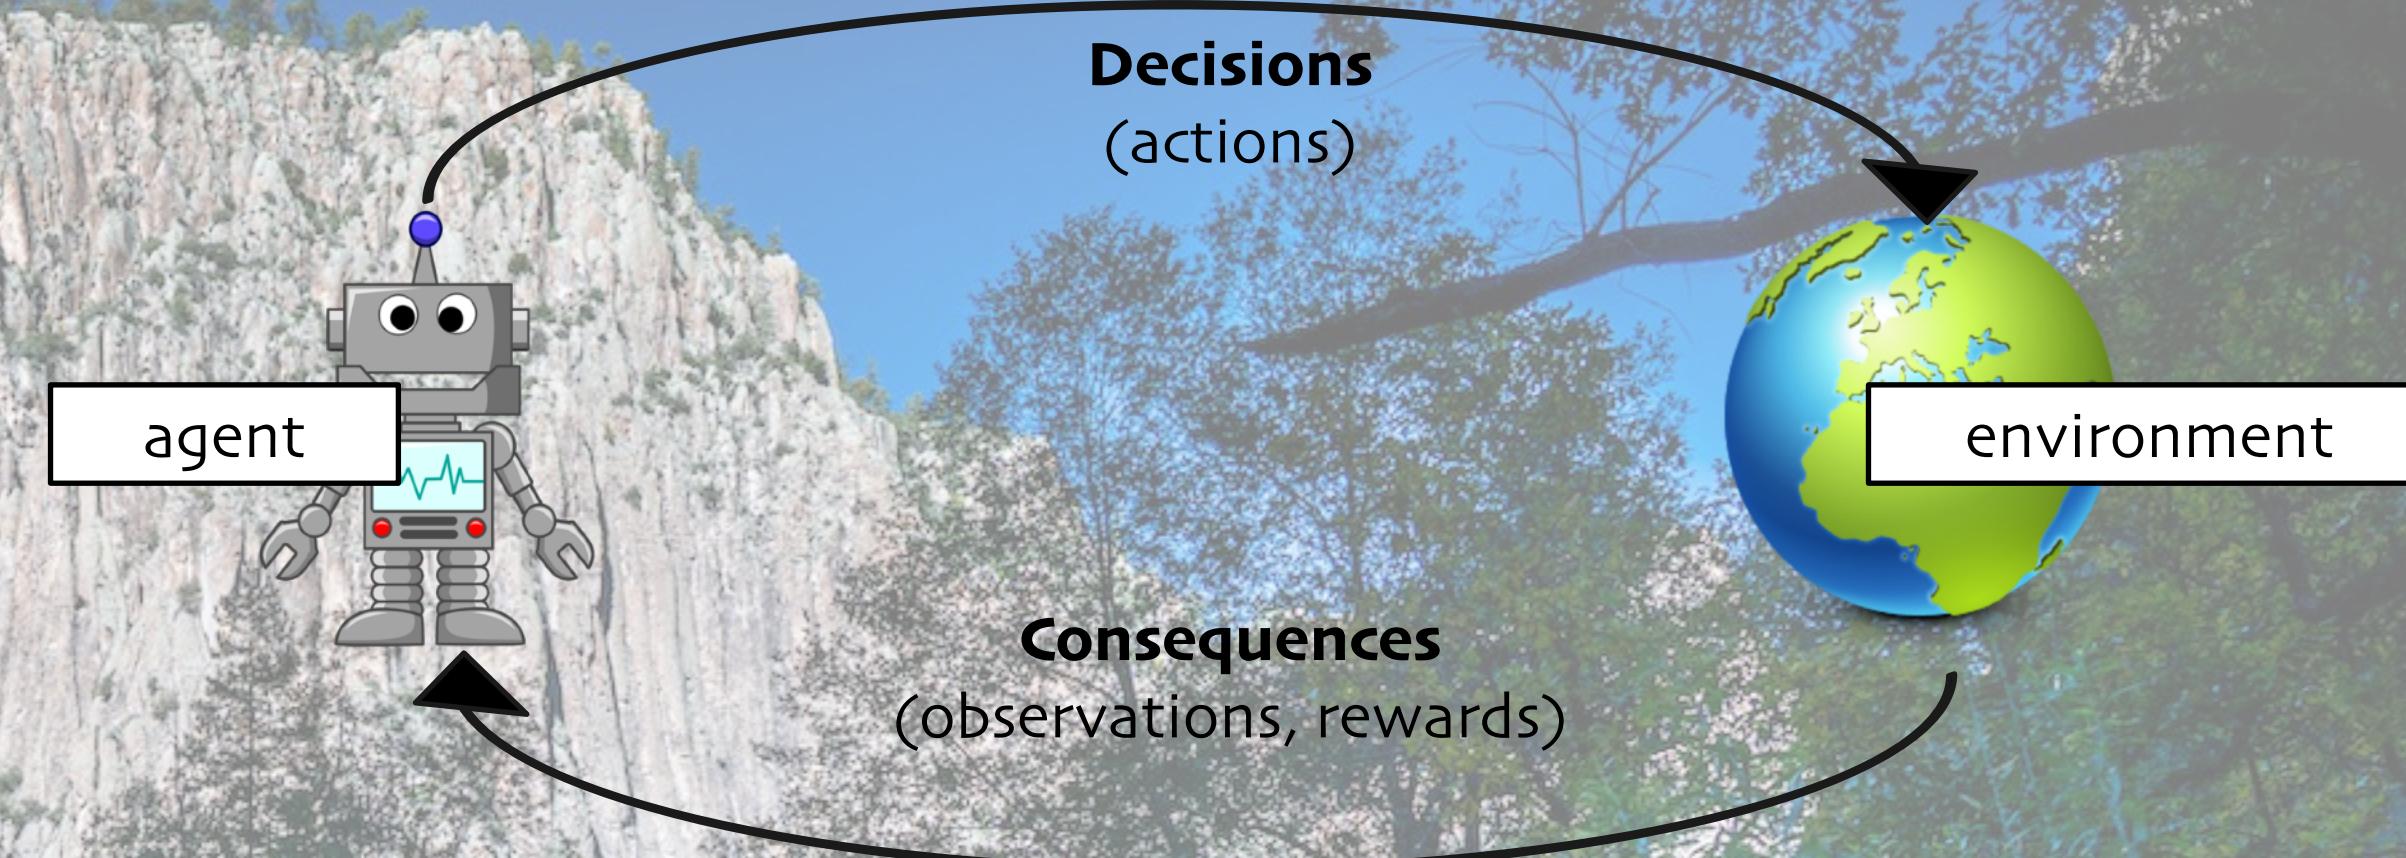
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

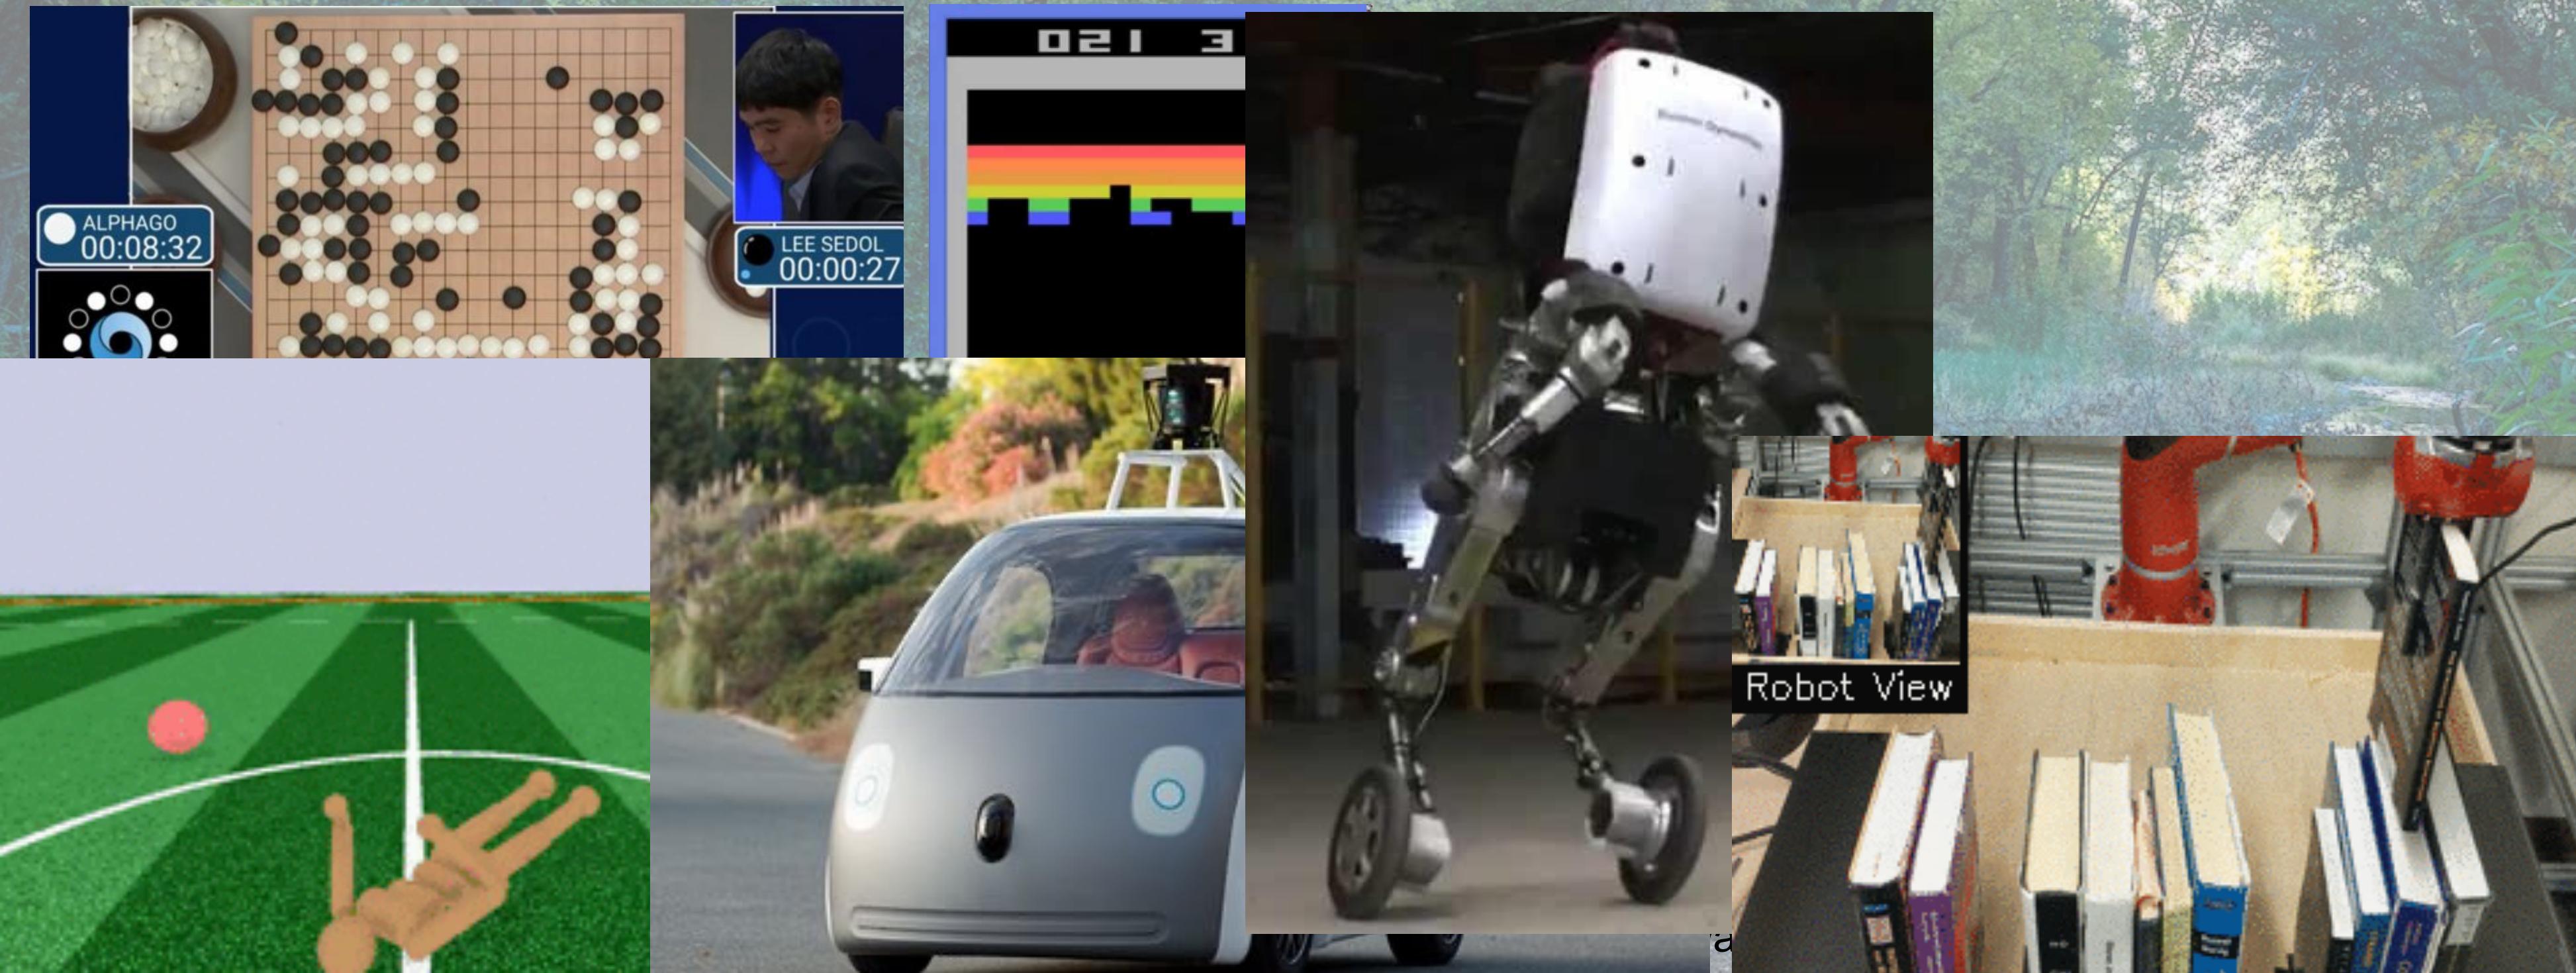
Industrial
Processes

System
Optimization

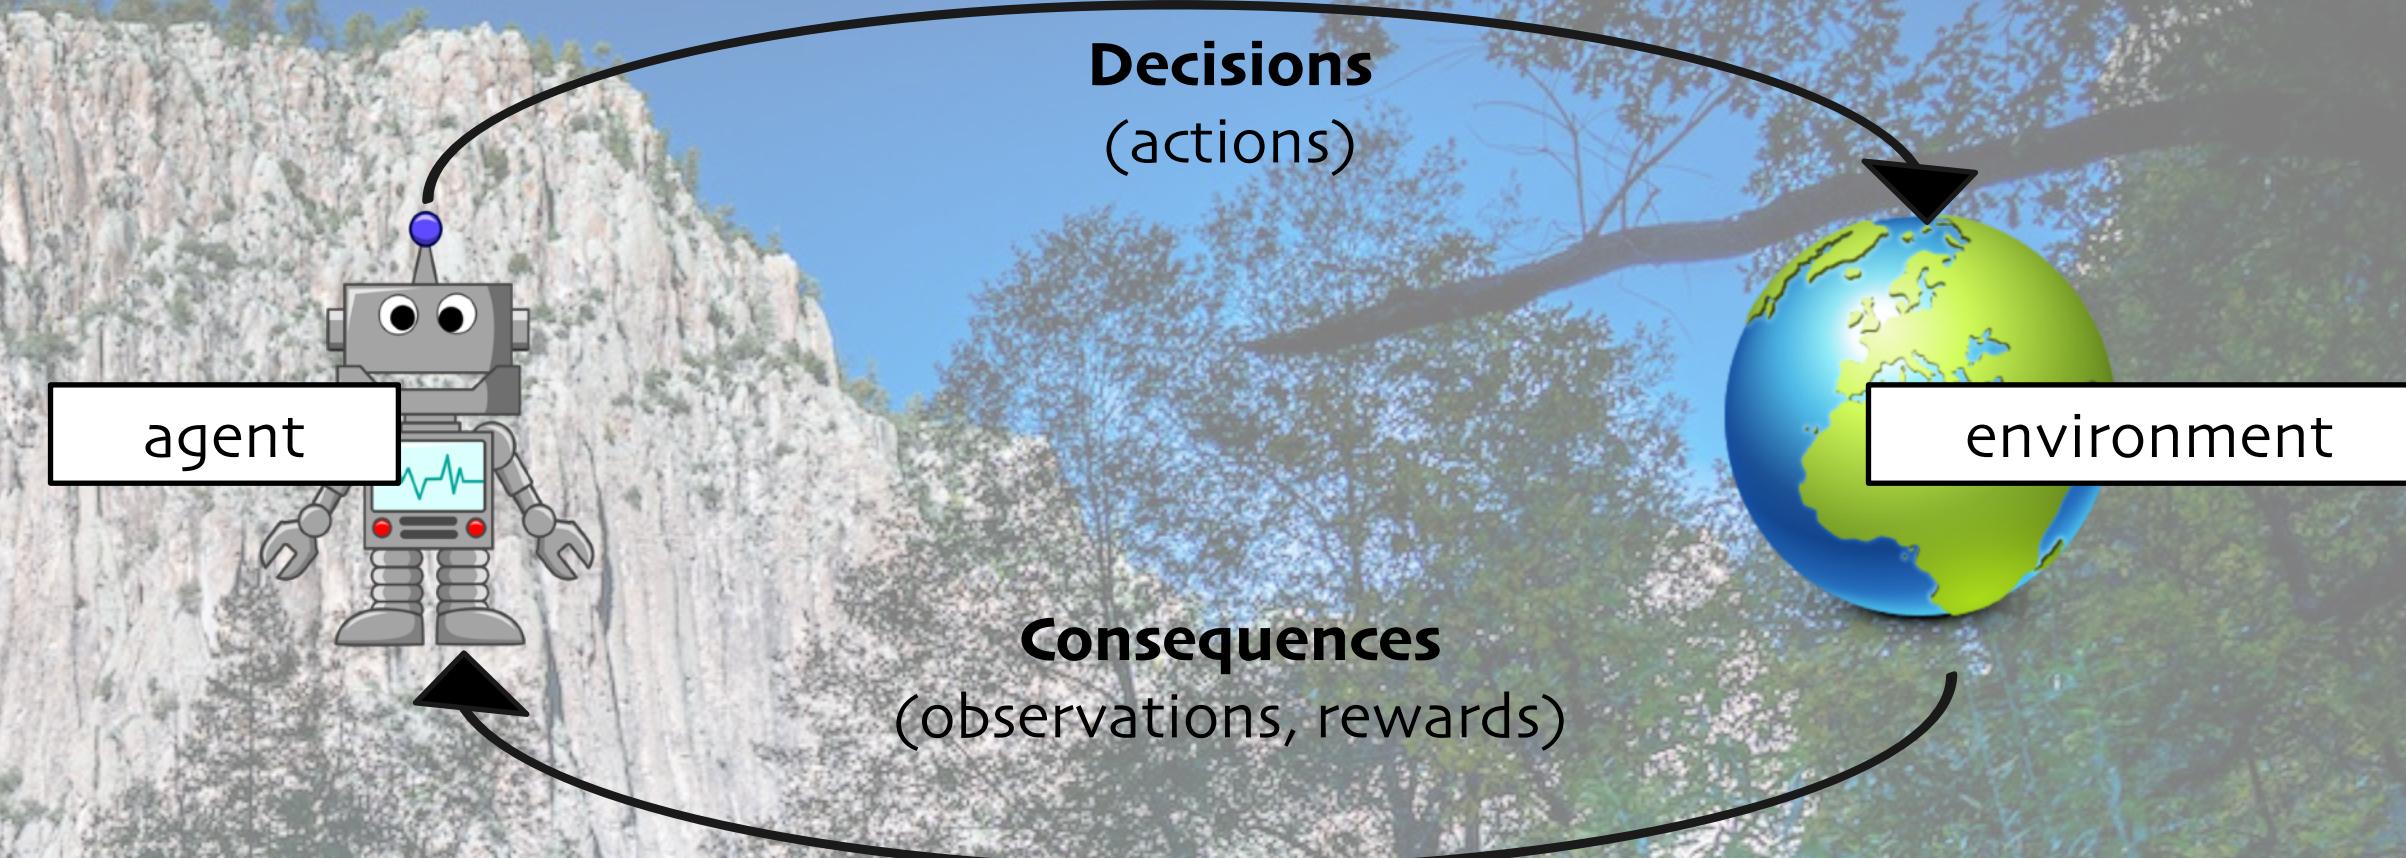
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

Industrial
Processes

System
Optimization

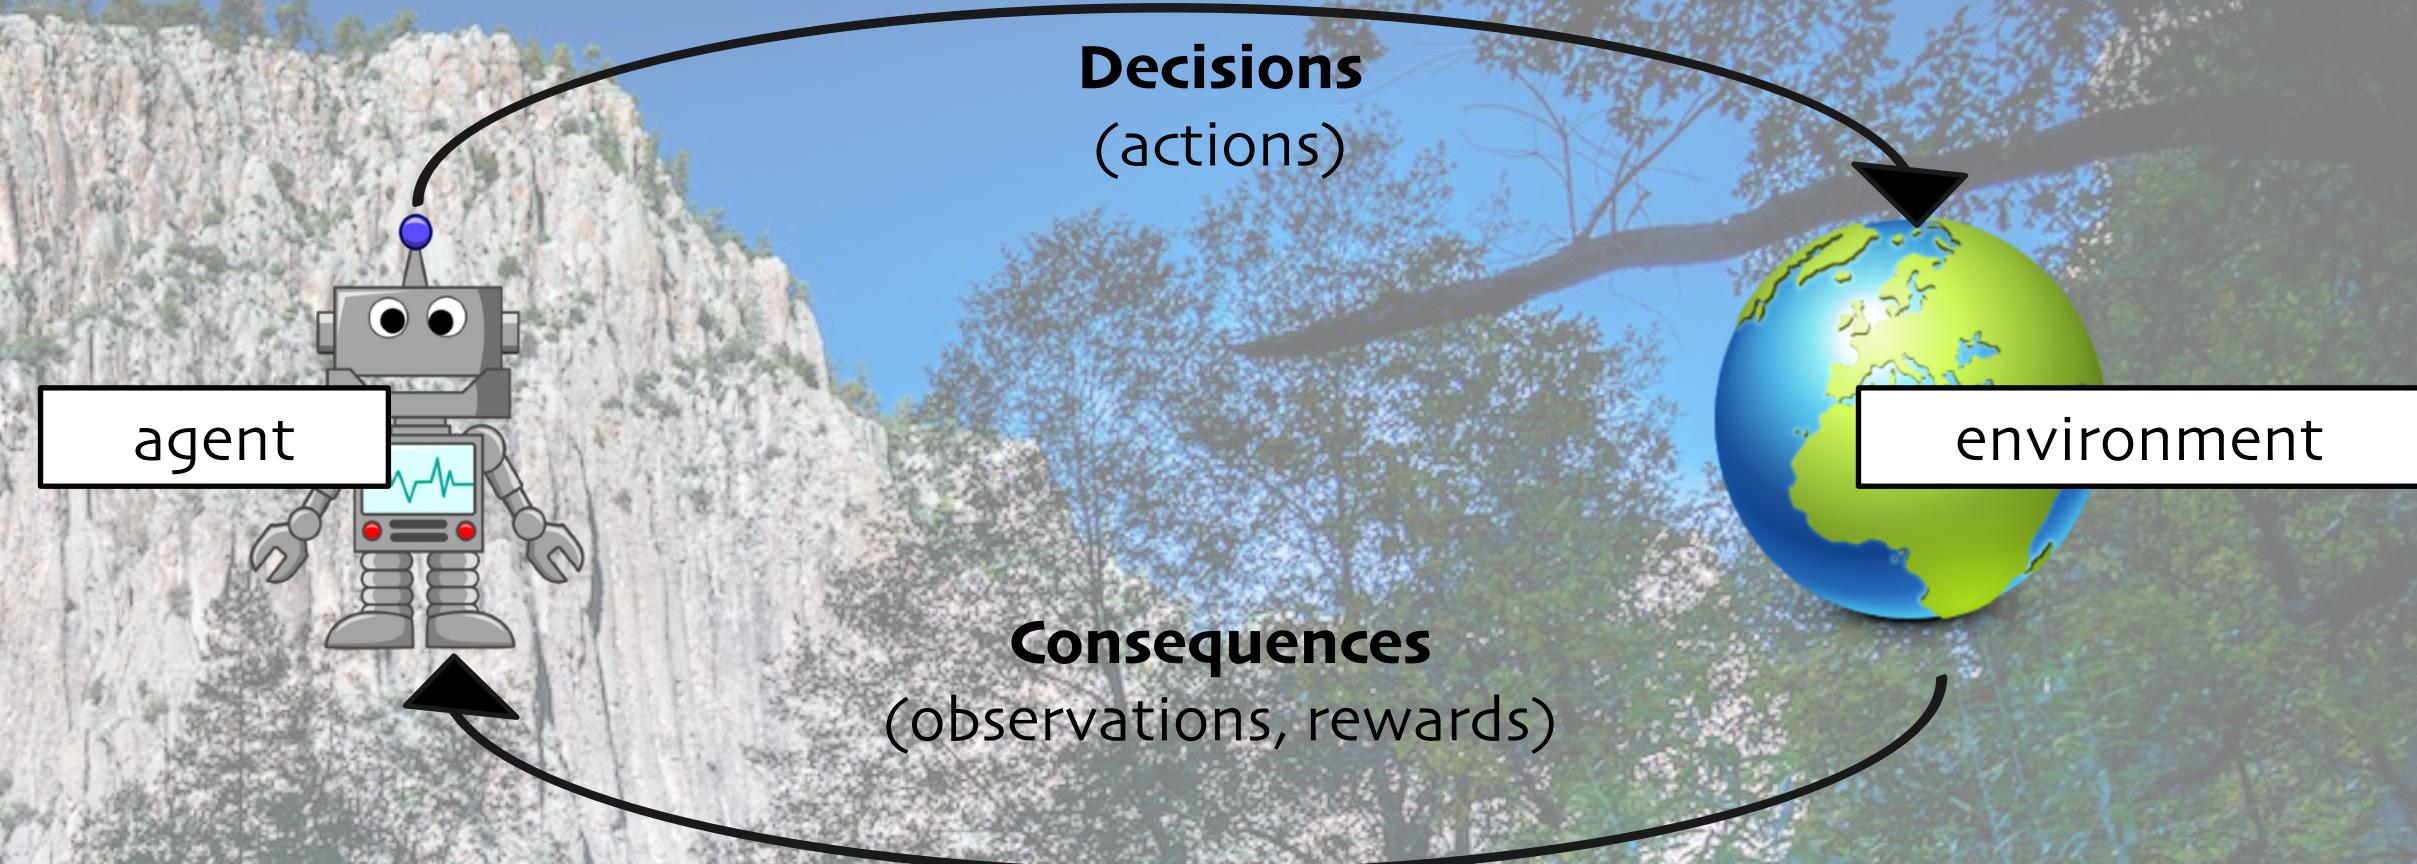
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

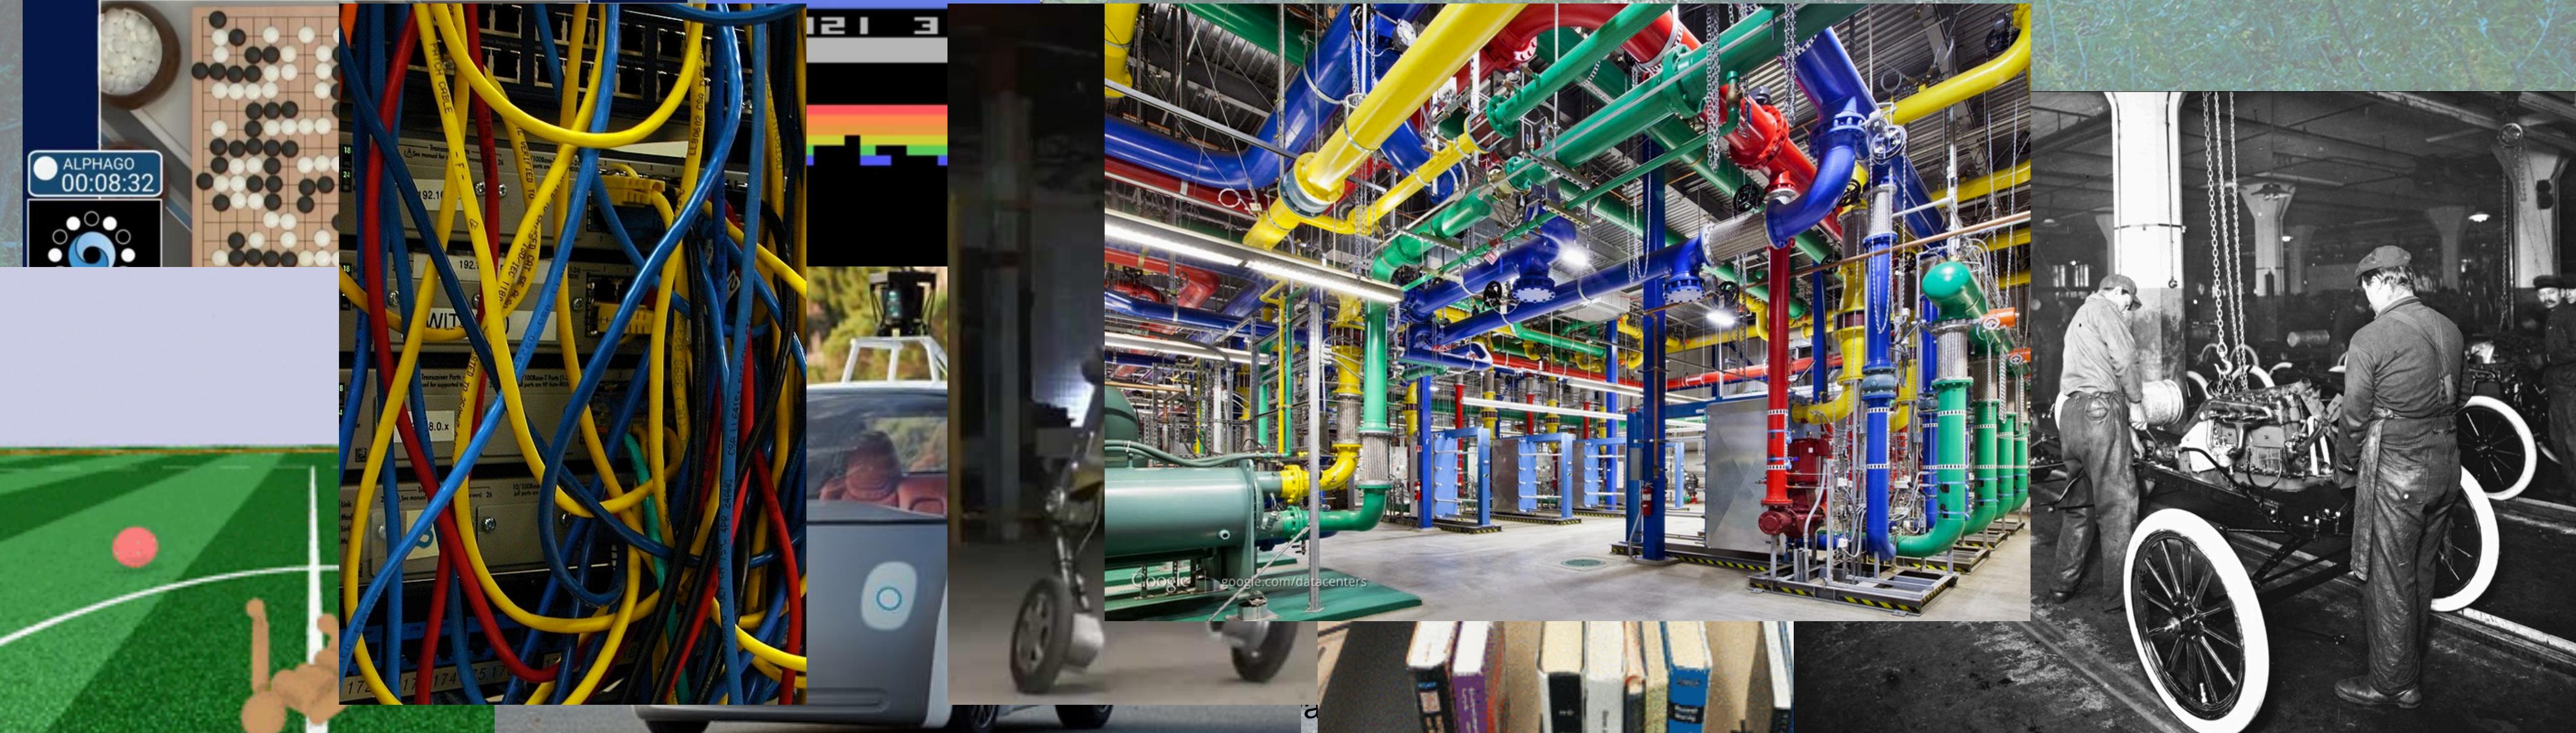
Industrial
Processes

System
Optimization

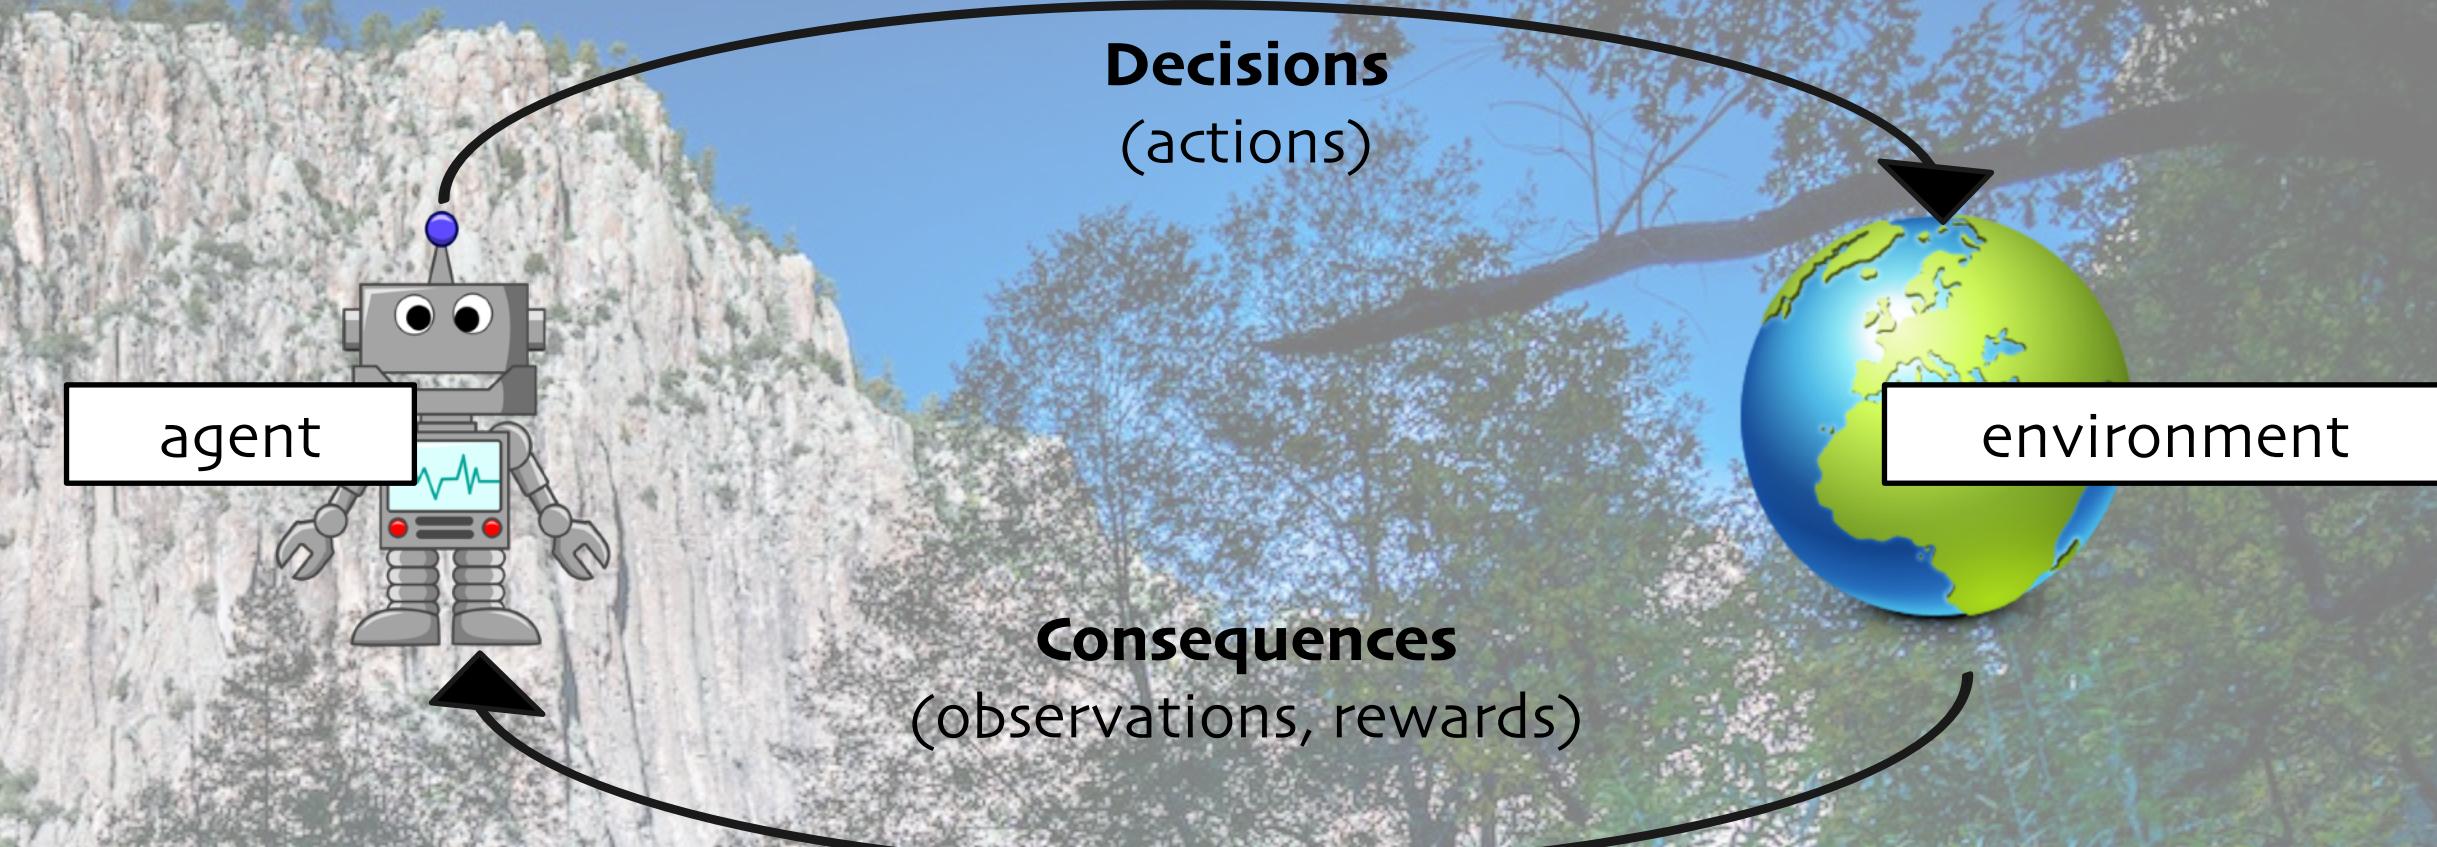
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

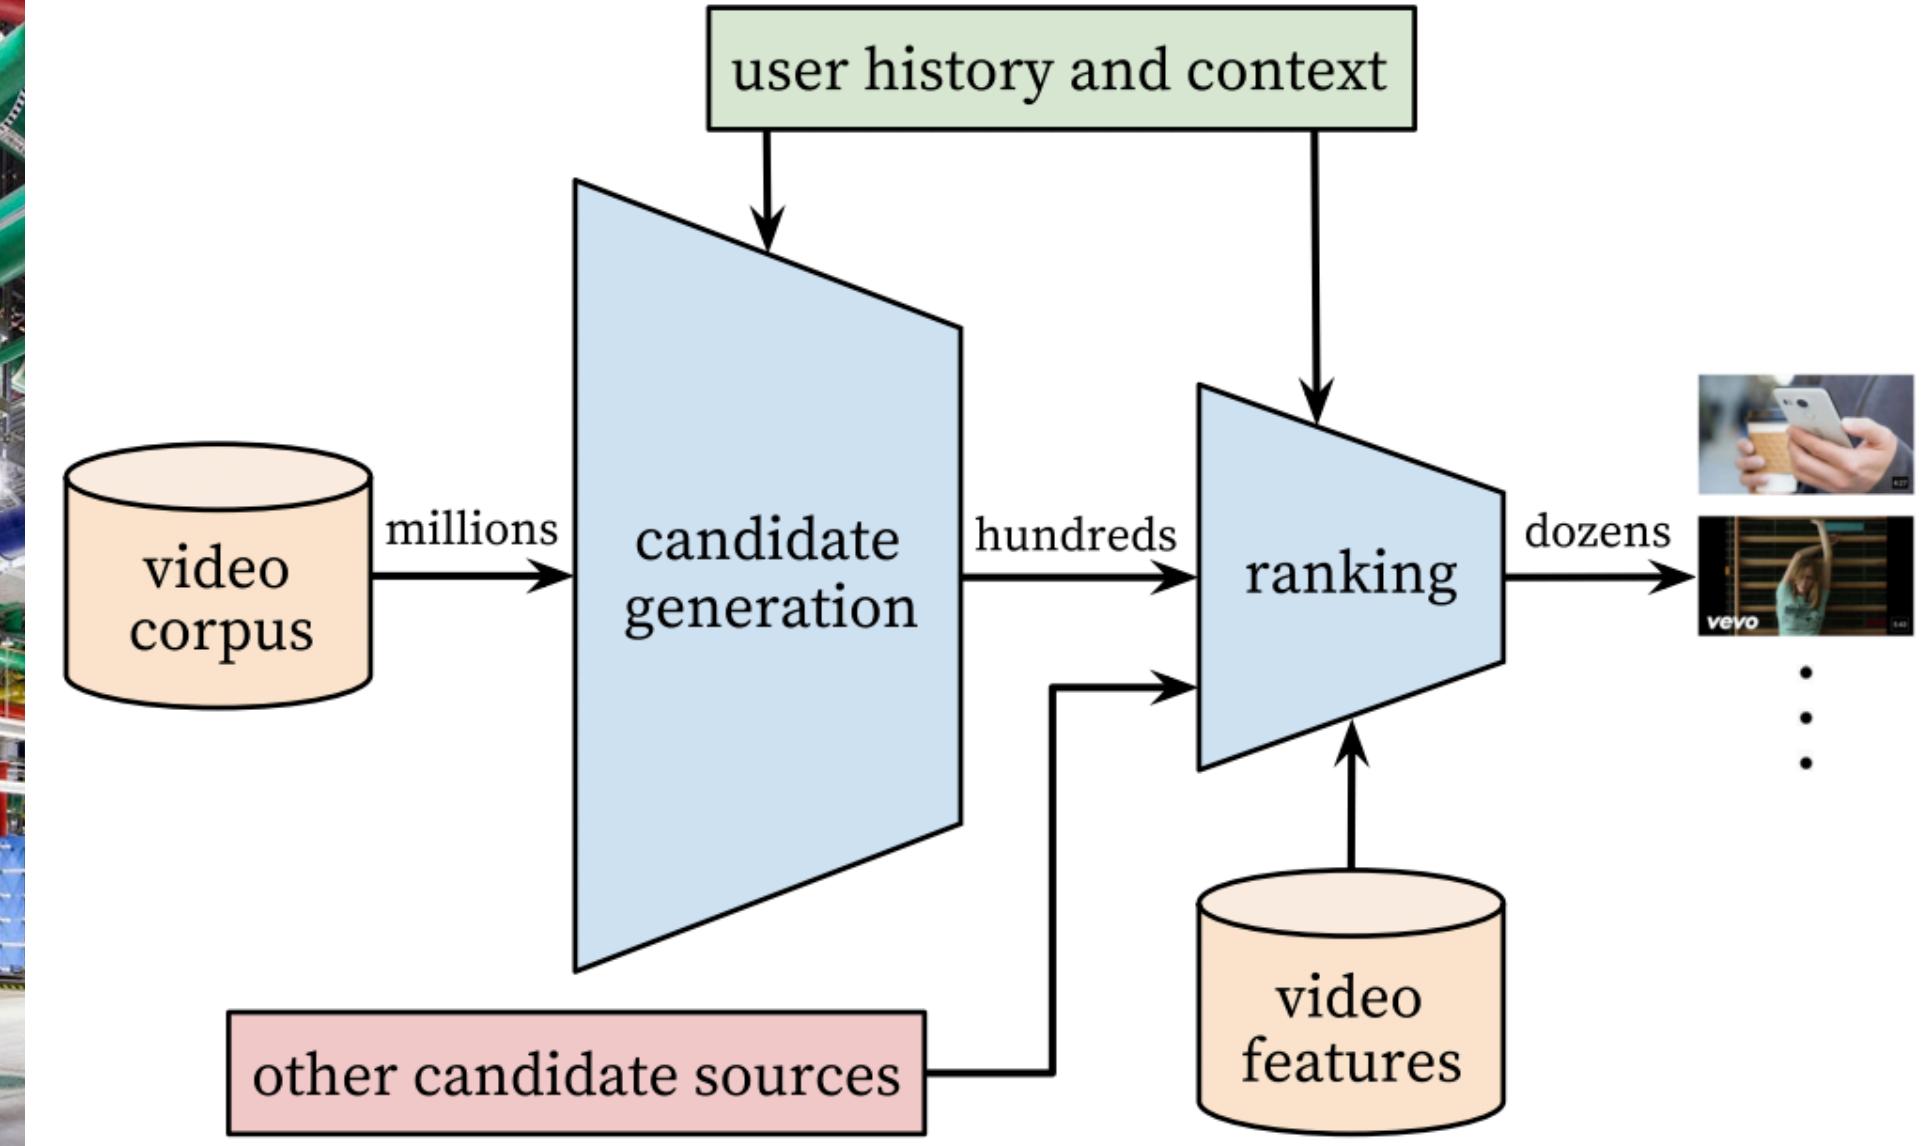
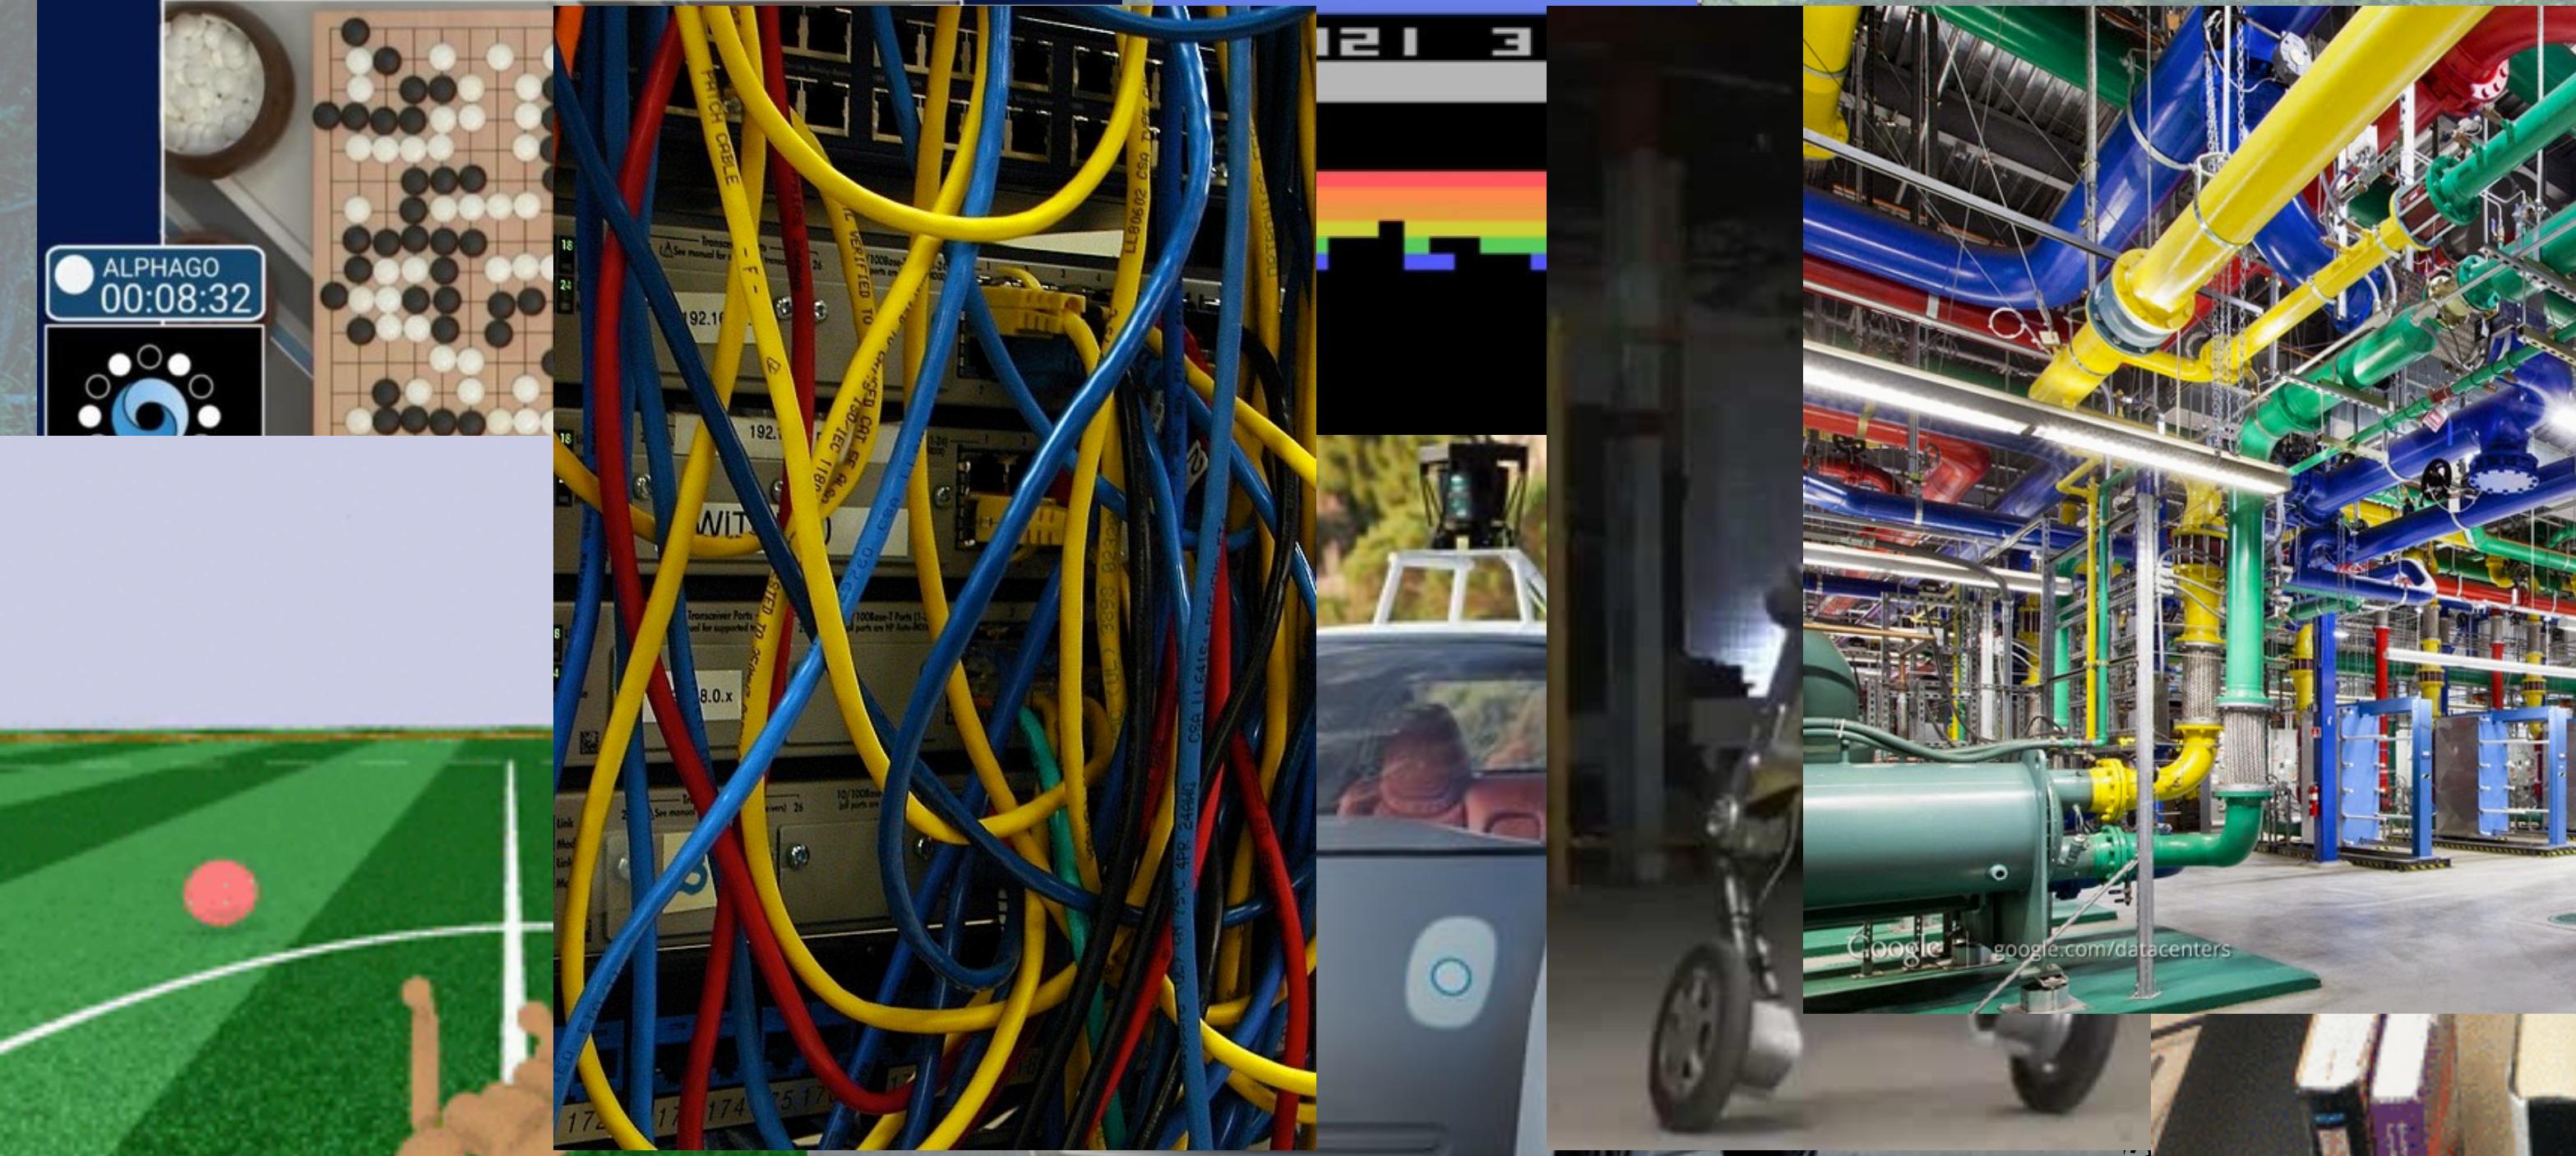
Industrial
Processes

System
Optimization

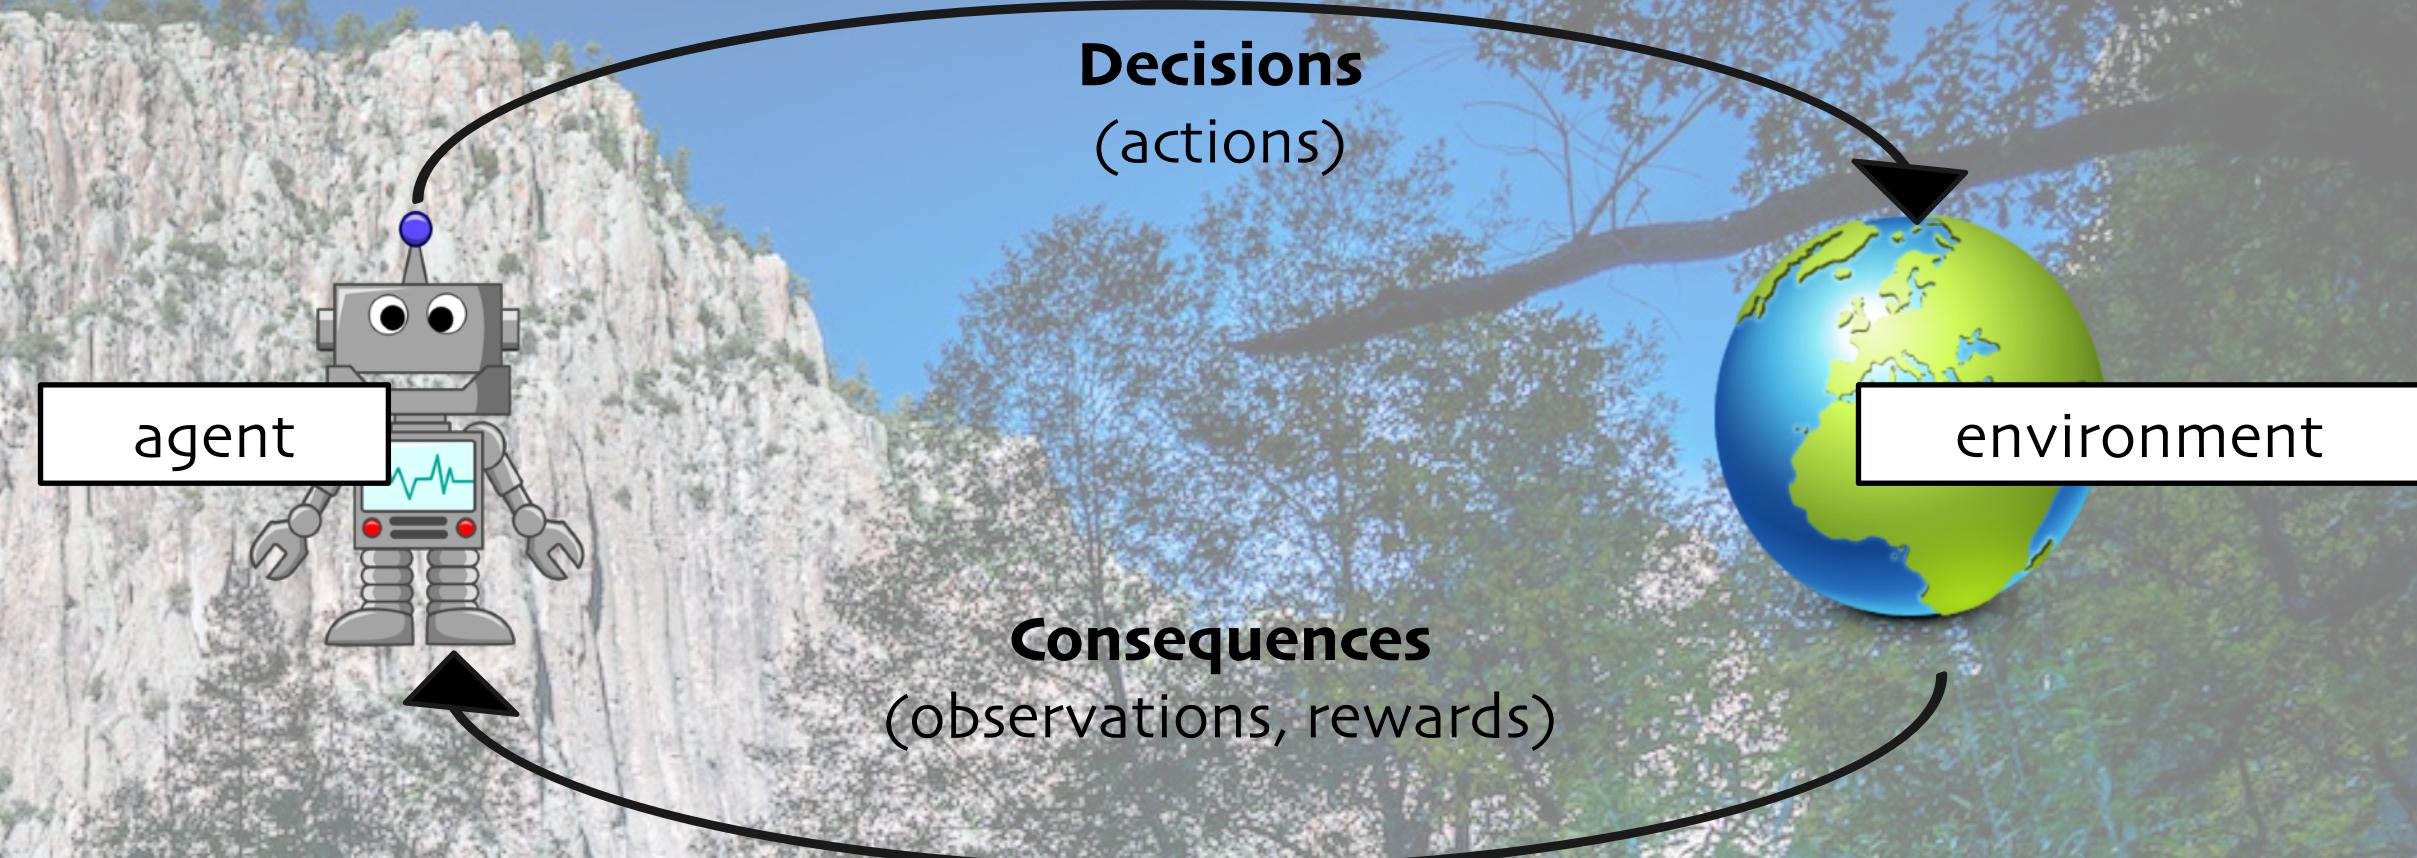
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

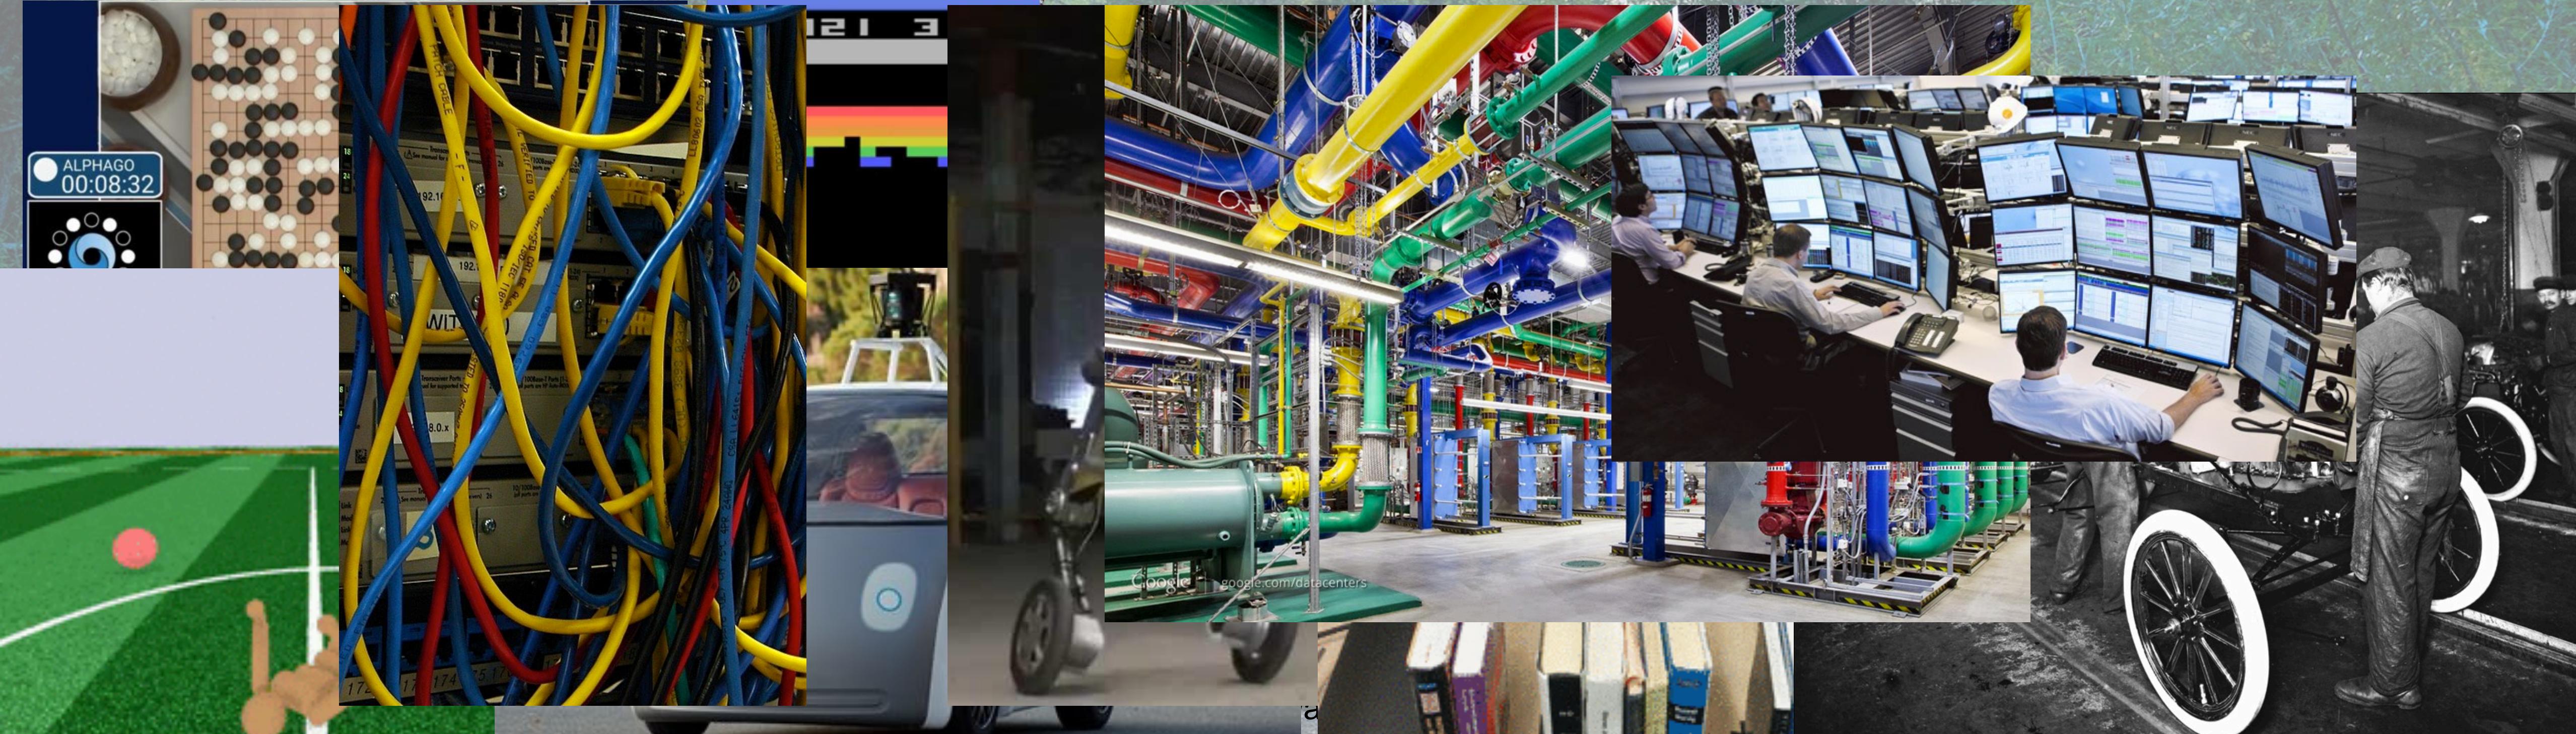
Industrial
Processes

System
Optimization

Advertising,
Recommendations

Finance

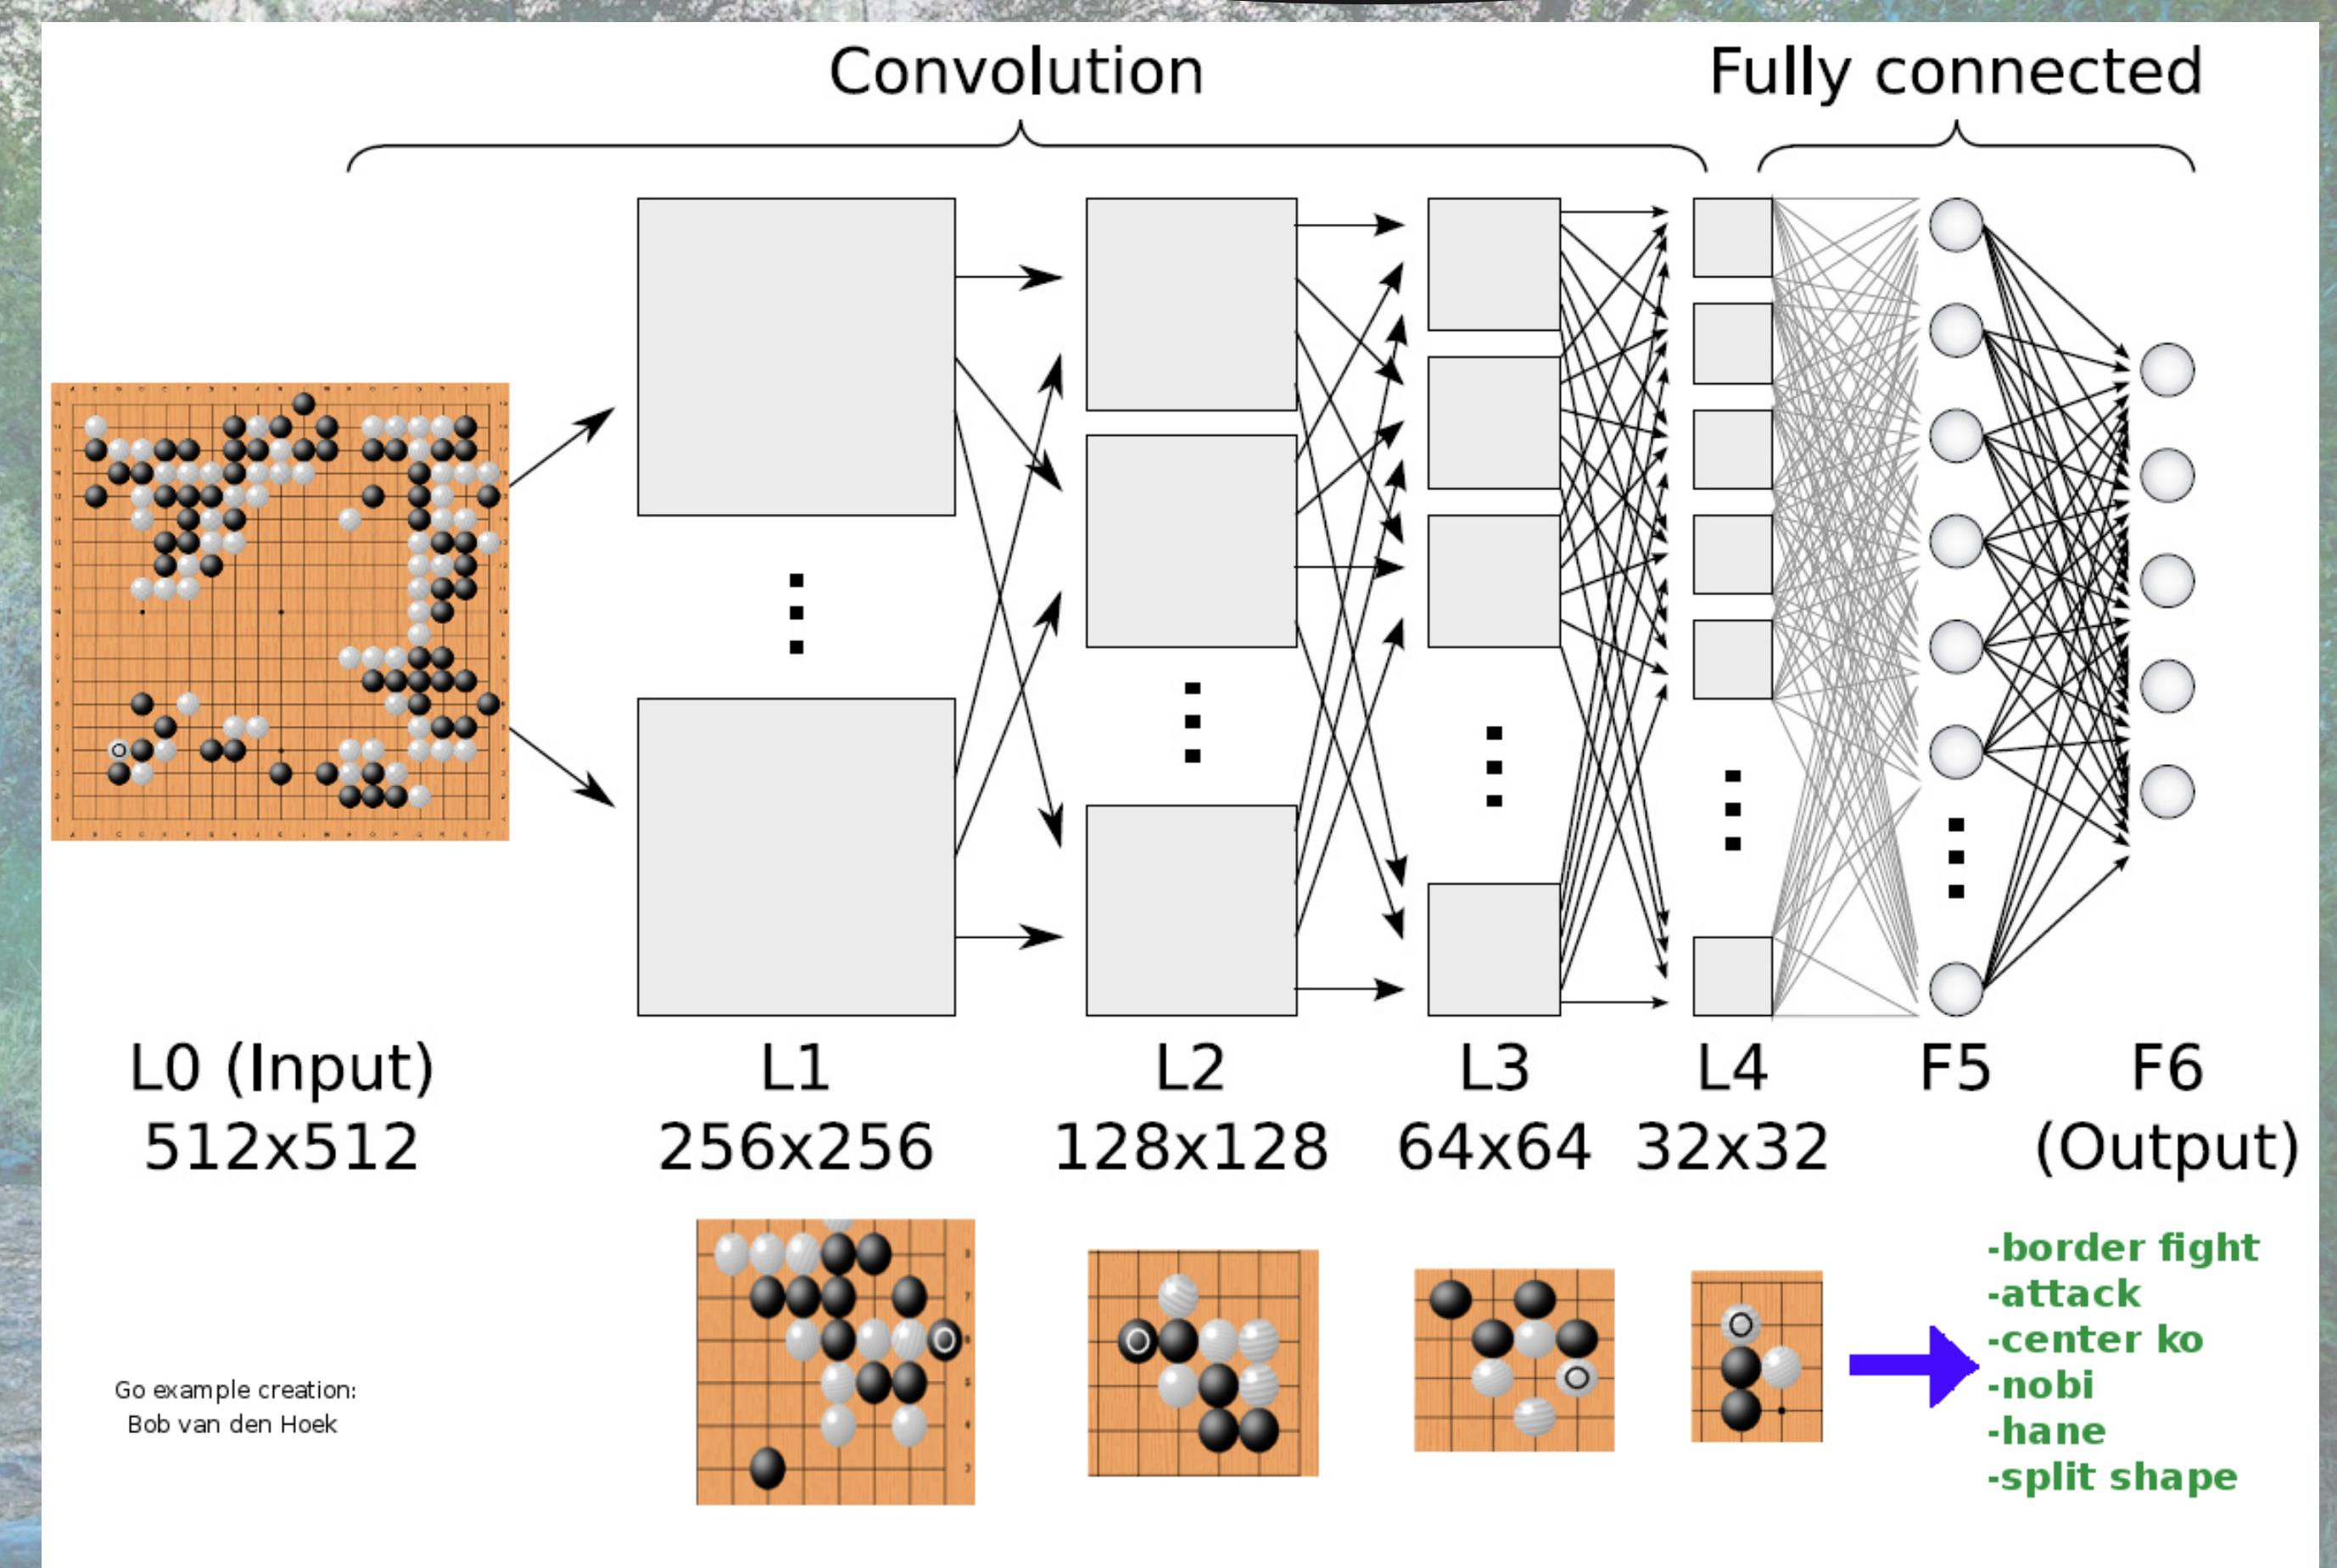
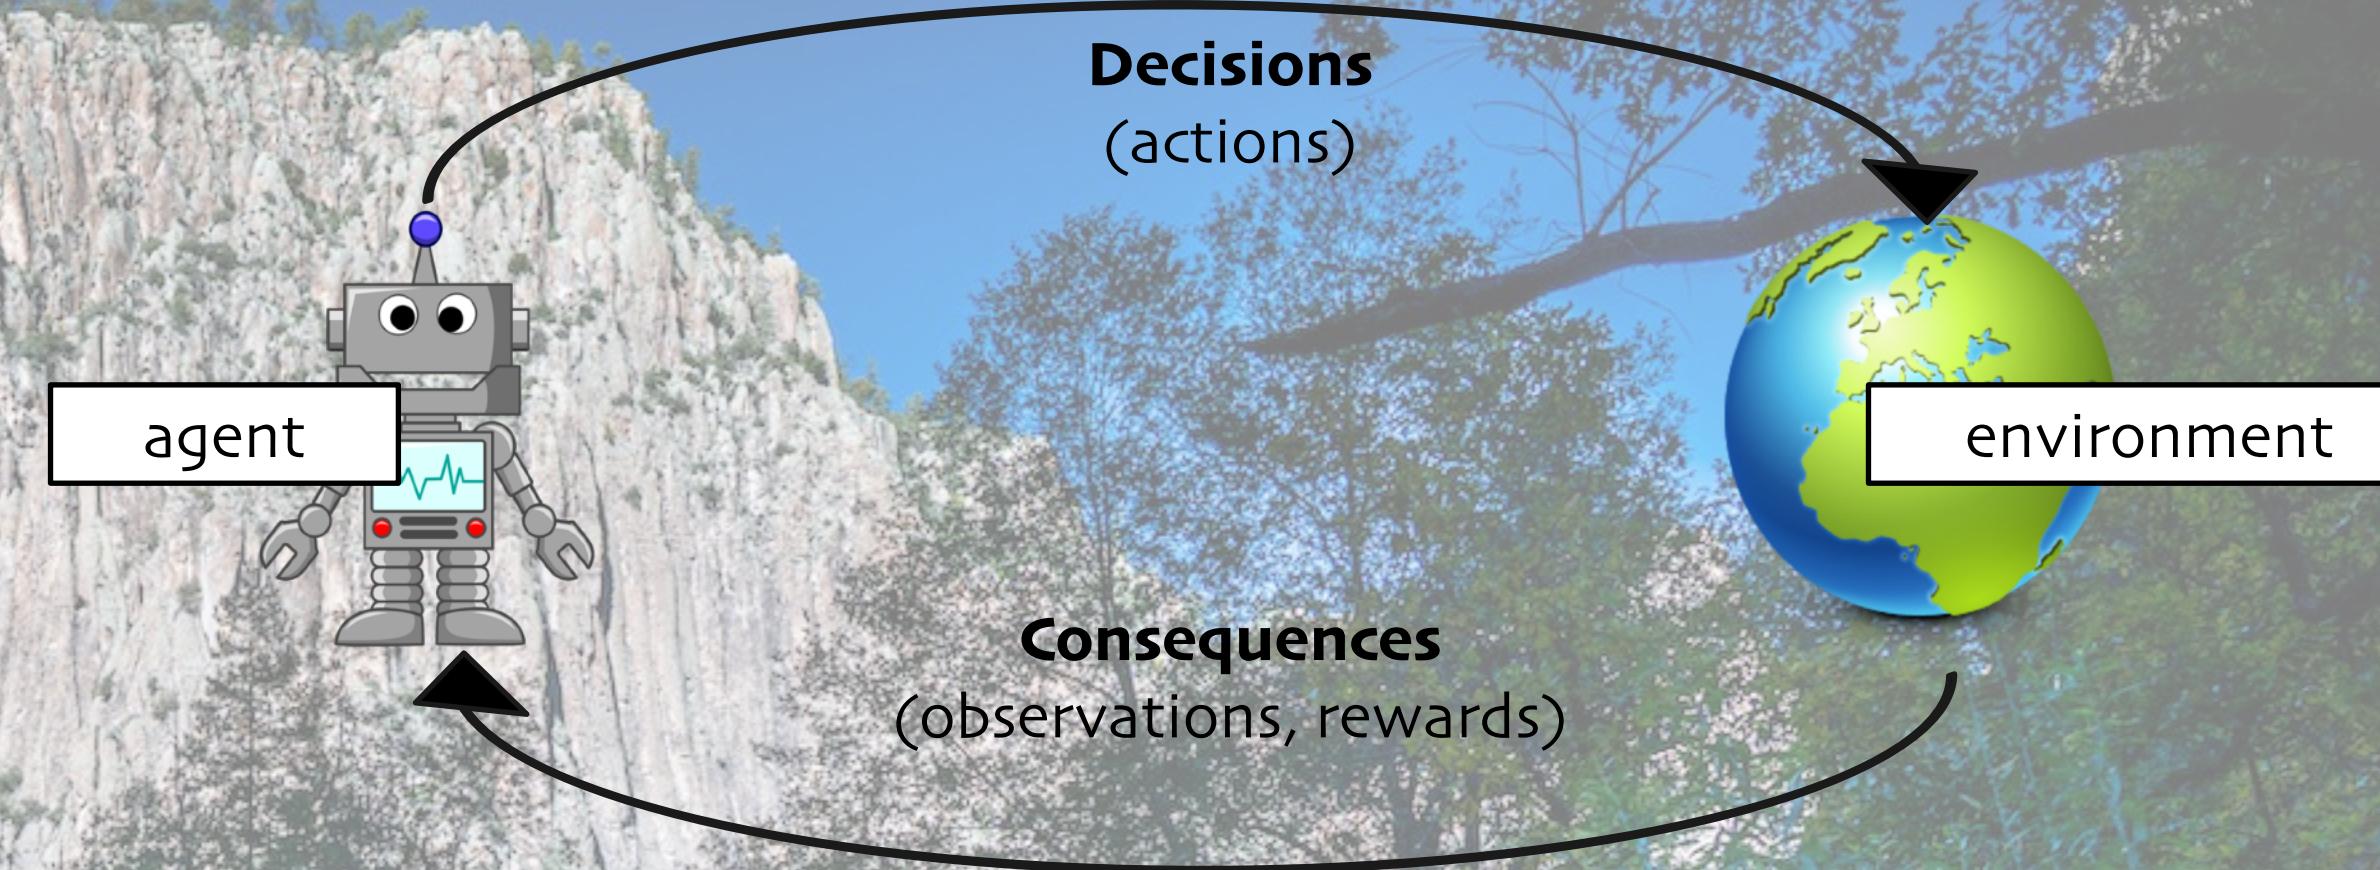
RL applications



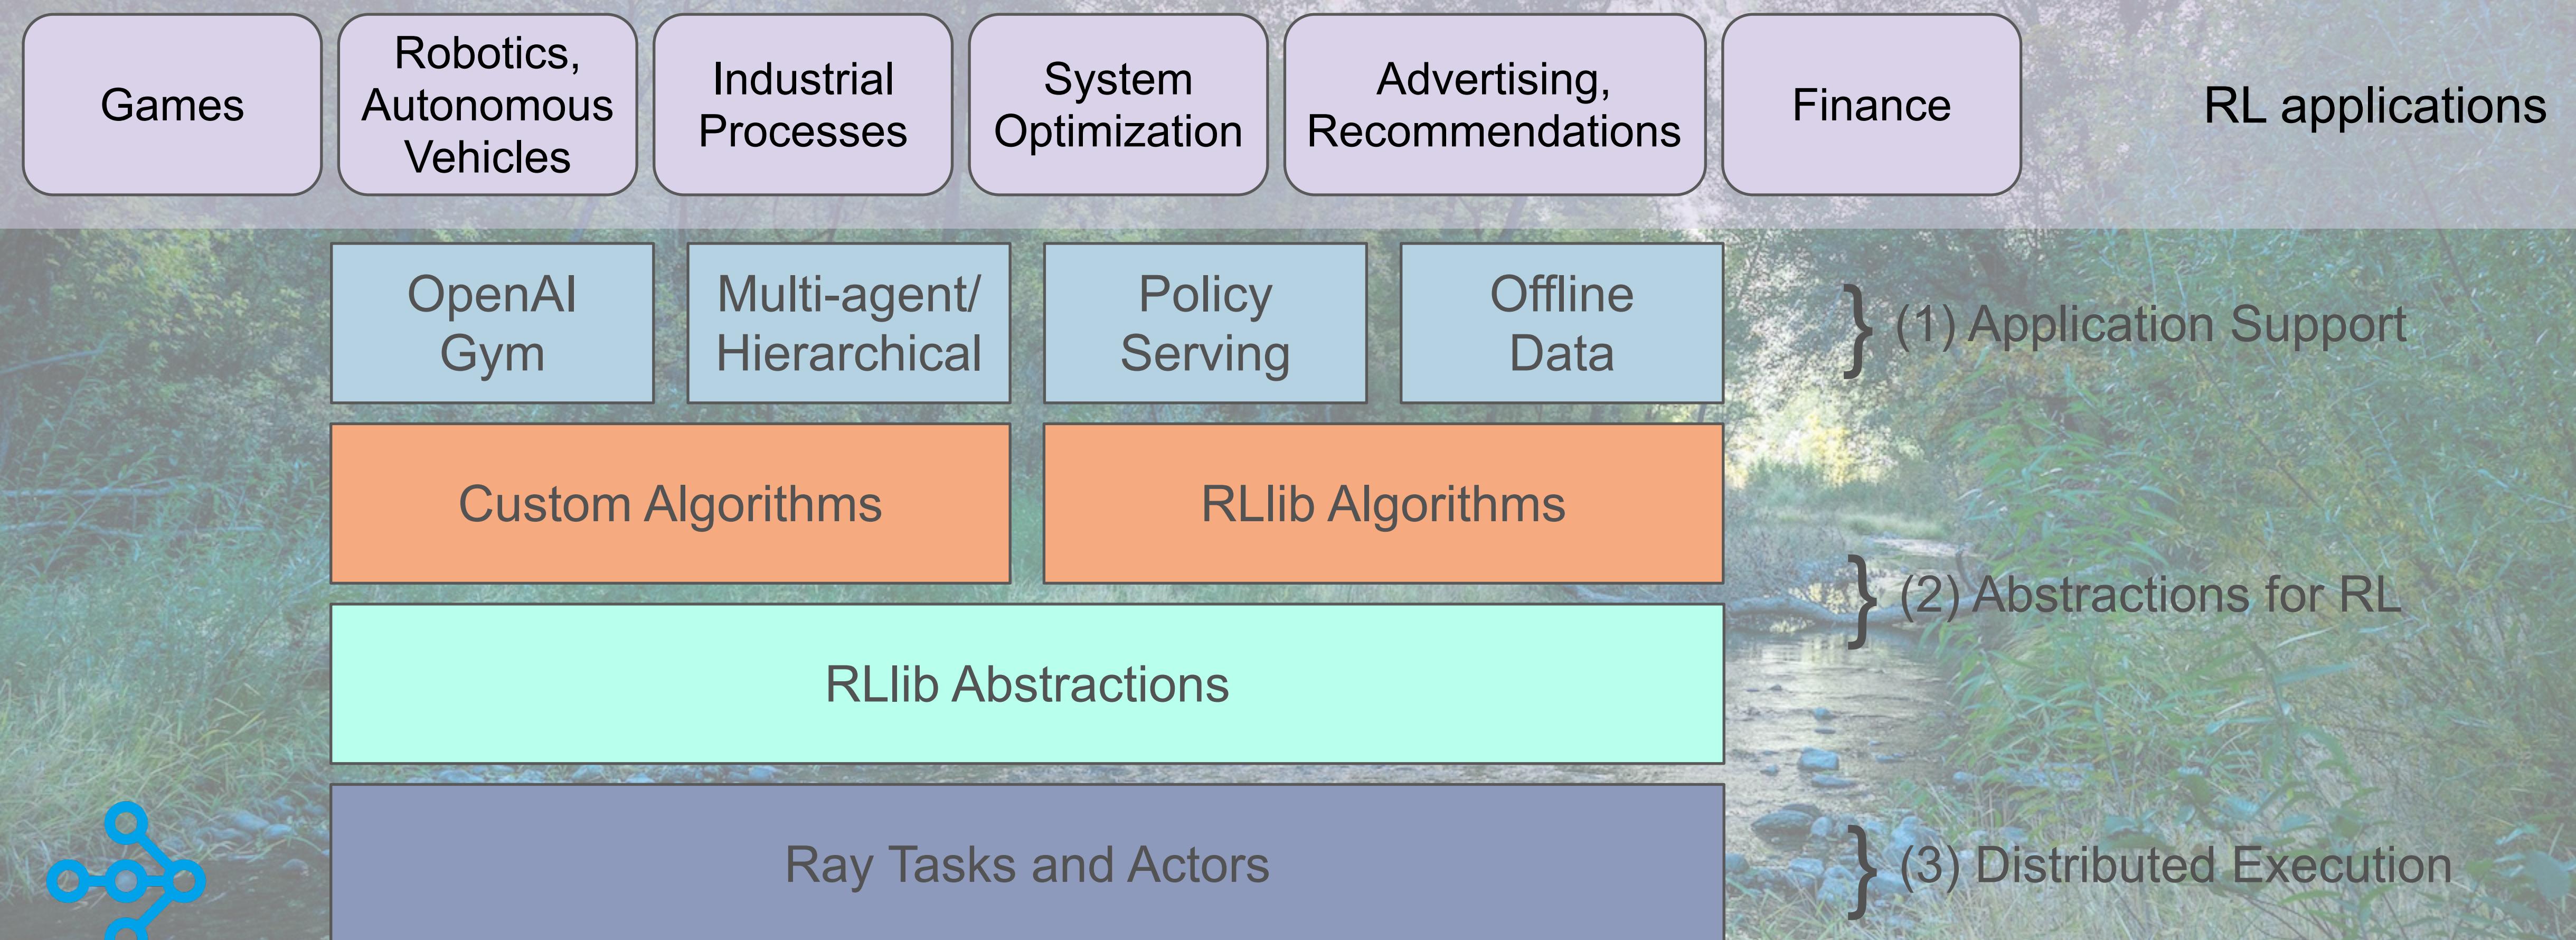
Go as a Reinforcement Learning Problem

AlphaGo (Silver et al. 2016)

- **Observations:**
 - board state
- **Actions:**
 - where to place the stones
- **Rewards:**
 - 1 if win
 - 0 otherwise



RLLib: A Scalable, Unified Library for RL



A Broad Range of Popular Algorithms

- High-throughput architectures
 - [Distributed Prioritized Experience Replay \(Ape-X\)](#)
 - [Importance Weighted Actor-Learner Architecture \(IMPALA\)](#)
 - [Asynchronous Proximal Policy Optimization \(APPO\)](#)
- Gradient-based
 - [Soft Actor-Critic \(SAC\)](#)
 - [Advantage Actor-Critic \(A₂C, A₃C\)](#)
 - [Deep Deterministic Policy Gradients \(DDPG, TD3\)](#)
 - [Deep Q Networks \(DQN, Rainbow, Parametric DQN\)](#)
 - [Policy Gradients](#)
 - [Proximal Policy Optimization \(PPO\)](#)
- gradient-free
 - [Augmented Random Search \(ARS\)](#)
 - [Evolution Strategies](#)
- Multi-agent specific
 - [QMIX Monotonic Value Factorisation \(QMIX, VDN, IQN\)](#)
- Offline
 - [Advantage Re-Weighted Imitation Learning \(MARWIL\)](#)



Amazon SageMaker RL

Reinforcement learning for every developer and data scientist



Amazon SageMaker RL

End-to-end examples for classic RL and real-world RL applications

Robotics

Industrial Control

HVAC

Autonomous Vehicles

Operations

Finance

Games

NLP

RL Environments to model real-world problems

AWS Simulation Environments

Amazon Sumerian

AWS RoboMaker

Open Source Environments

EnergyPlus

RoboSchool

PyBullet

...

Custom Environments

Bring Your Own

Commercial simulators

MATLAB & Simulink

Open AI Gym

RL Toolkits that provide RL agent algorithm implementations

RL-Coach

DQN

PPO

HER

Rainbow

...

RL-Ray RLLib

APEX

ES

IMPALA

A3C

...

Open AI Baselines

TRPO

GAIL

...

...

TensorFlow

MxNet

PyTorch

Chainer

Training Options

Single Machine / Distributed

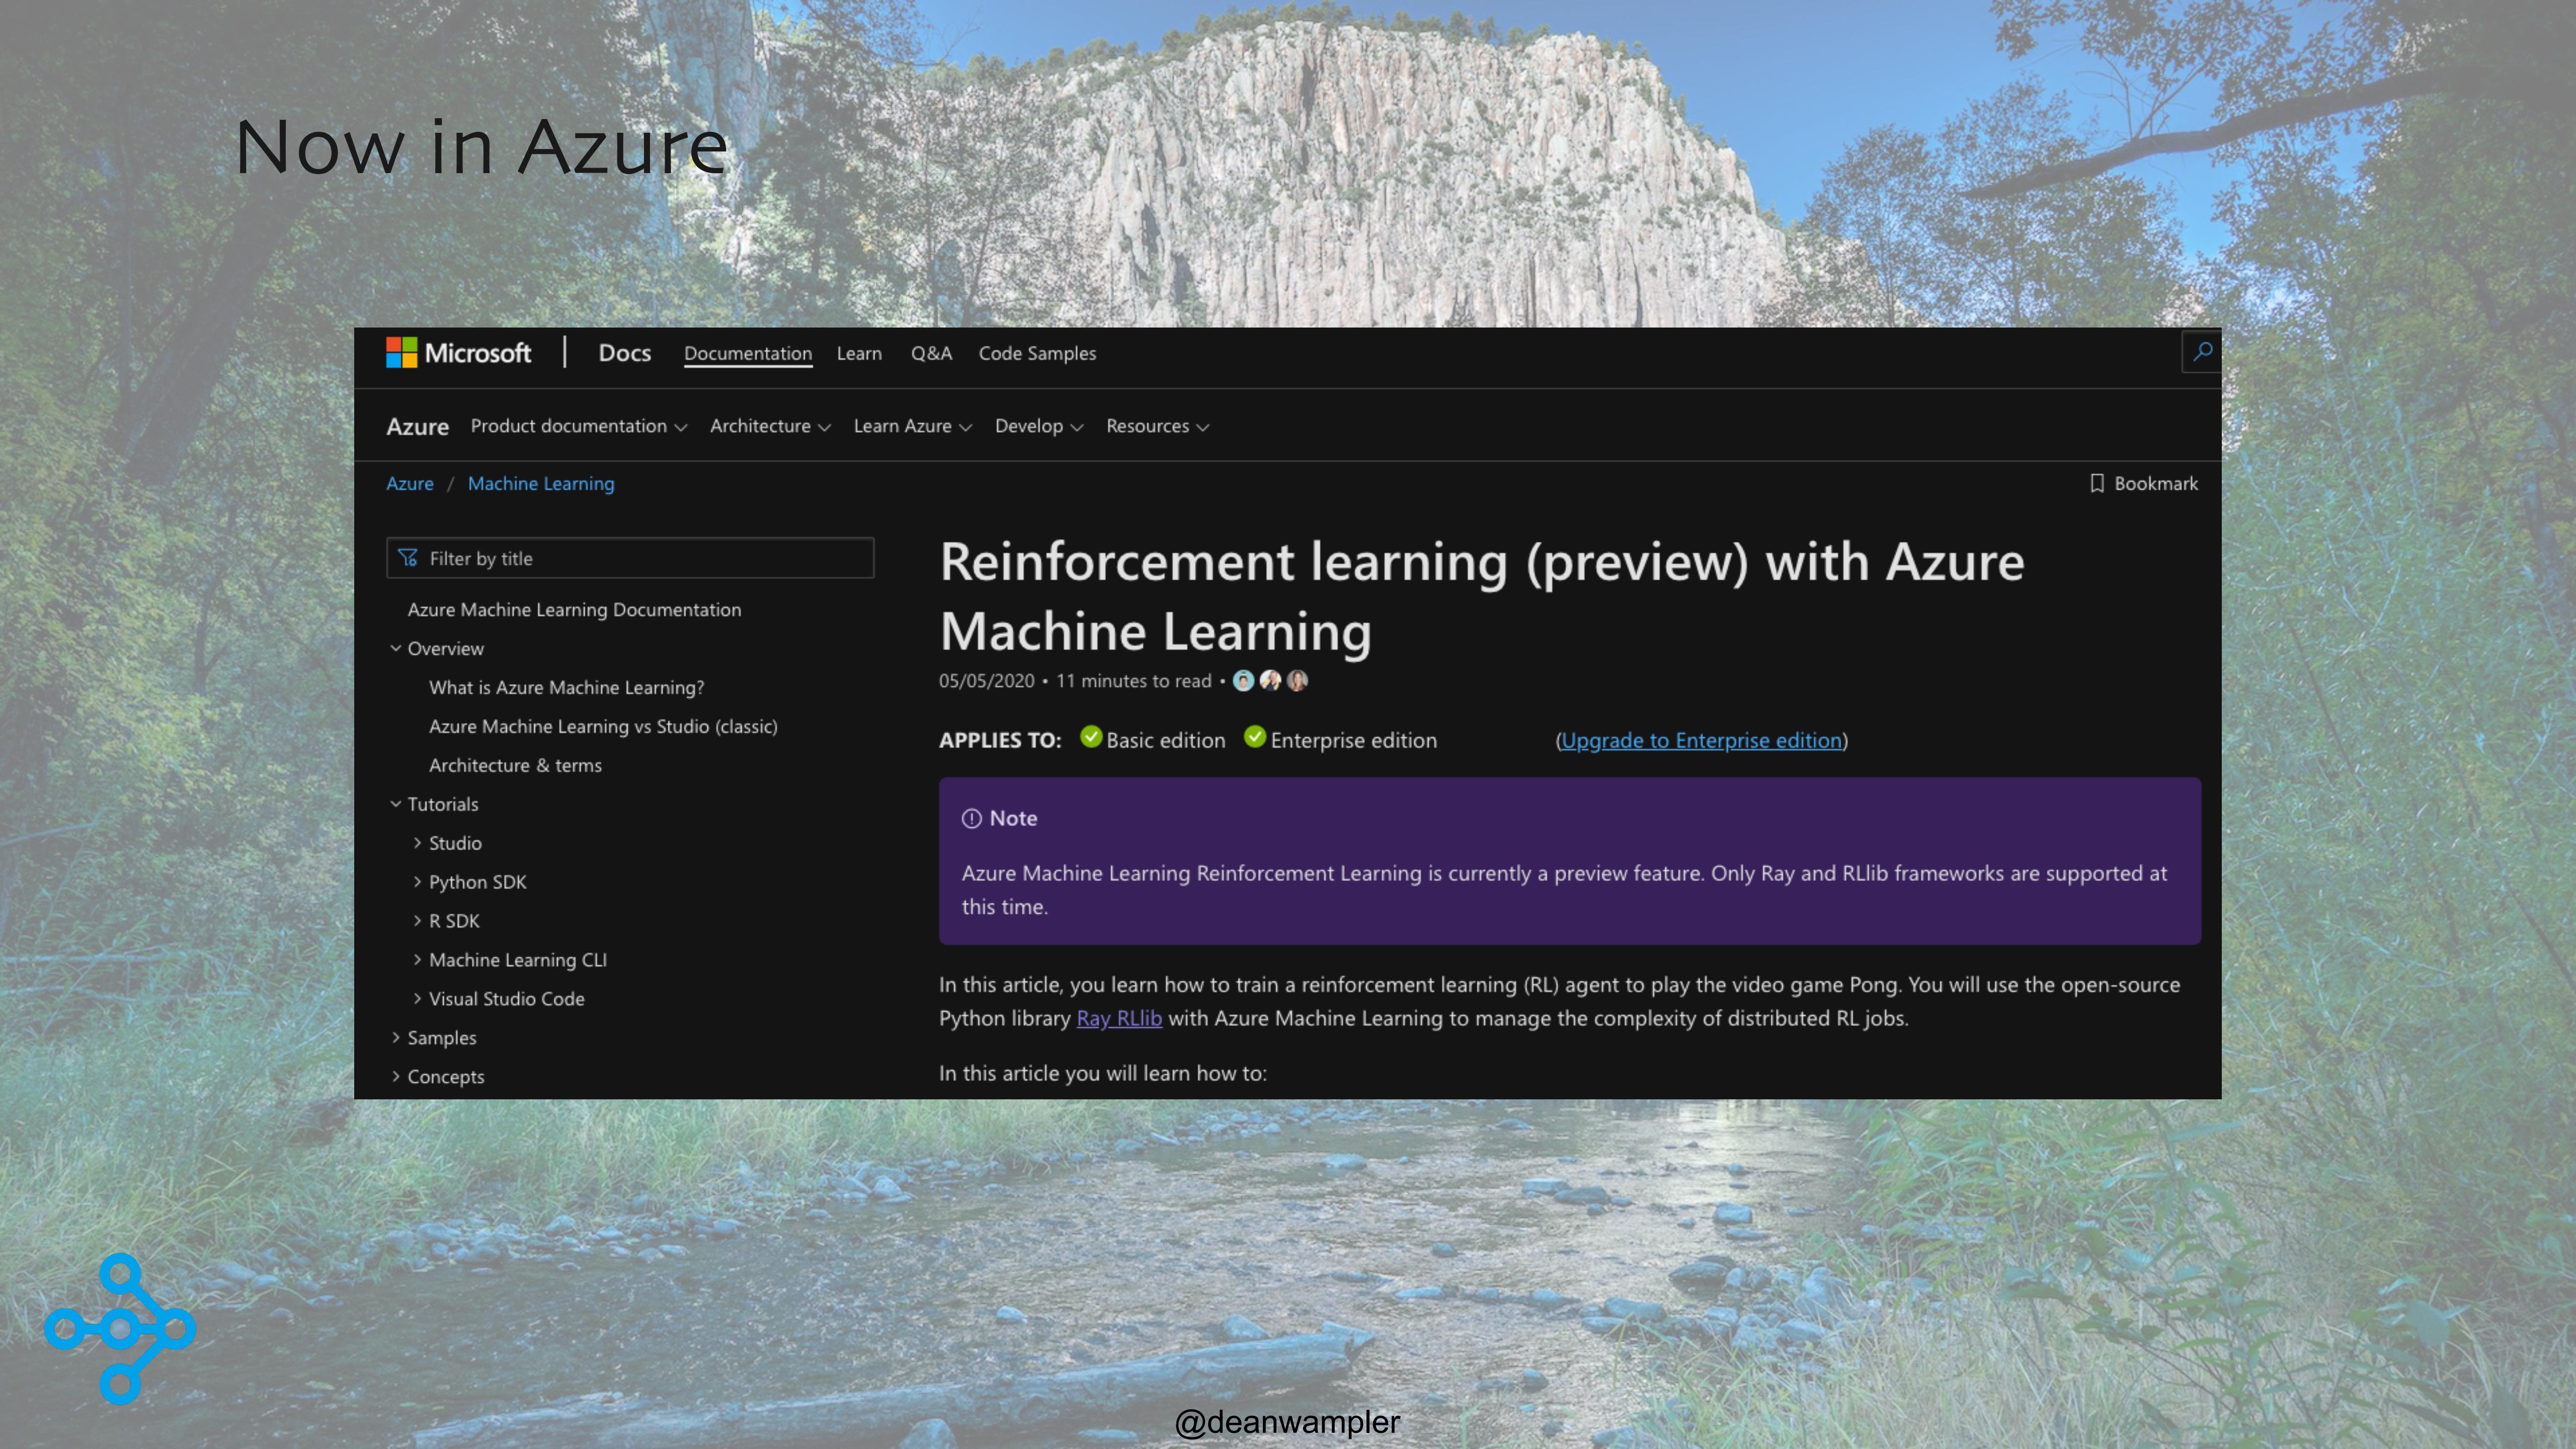
Local / Remote simulation

CPU / GPU Hardware

SageMaker supported

Customer BYO

Now in Azure



A Microsoft Docs page for "Reinforcement learning (preview) with Azure Machine Learning".

Microsoft | Docs Documentation Learn Q&A Code Samples

Azure Product documentation ▾ Architecture ▾ Learn Azure ▾ Develop ▾ Resources ▾

Azure / Machine Learning Bookmark

Filter by title

Azure Machine Learning Documentation

✓ Overview

- What is Azure Machine Learning?
- Azure Machine Learning vs Studio (classic)
- Architecture & terms

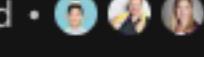
✓ Tutorials

- > Studio
- > Python SDK
- > R SDK
- > Machine Learning CLI
- > Visual Studio Code

> Samples

> Concepts

Reinforcement learning (preview) with Azure Machine Learning

05/05/2020 • 11 minutes to read • 

APPLIES TO:  Basic edition  Enterprise edition [\(Upgrade to Enterprise edition\)](#)

ⓘ Note

Azure Machine Learning Reinforcement Learning is currently a preview feature. Only Ray and RLlib frameworks are supported at this time.

In this article, you learn how to train a reinforcement learning (RL) agent to play the video game Pong. You will use the open-source Python library [Ray RLlib](#) with Azure Machine Learning to manage the complexity of distributed RL jobs.

In this article you will learn how to:

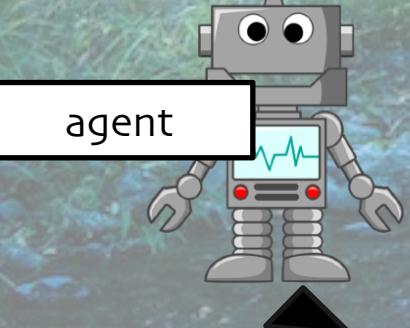


Diverse Compute Requirements Motivated Creation of Ray!

And repeated play,
over and over again,
to train for achieving
the best reward

Simulator (game
engine, robot sim,
factory floor sim...)

Complex agent?



Decisions (**actions**)

Consequences
(**observations, rewards**)

Neural network
“stuff”

L0 (Input)
512x512

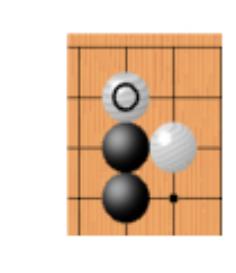
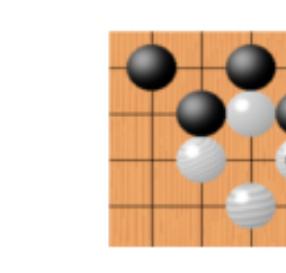
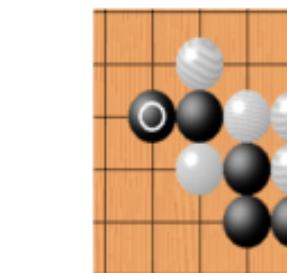
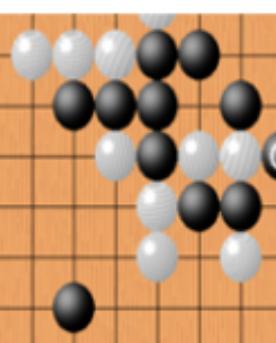
L1
256x256

L2
128x128

L3
64x64

L4
32x32

F5
F6
(Output)

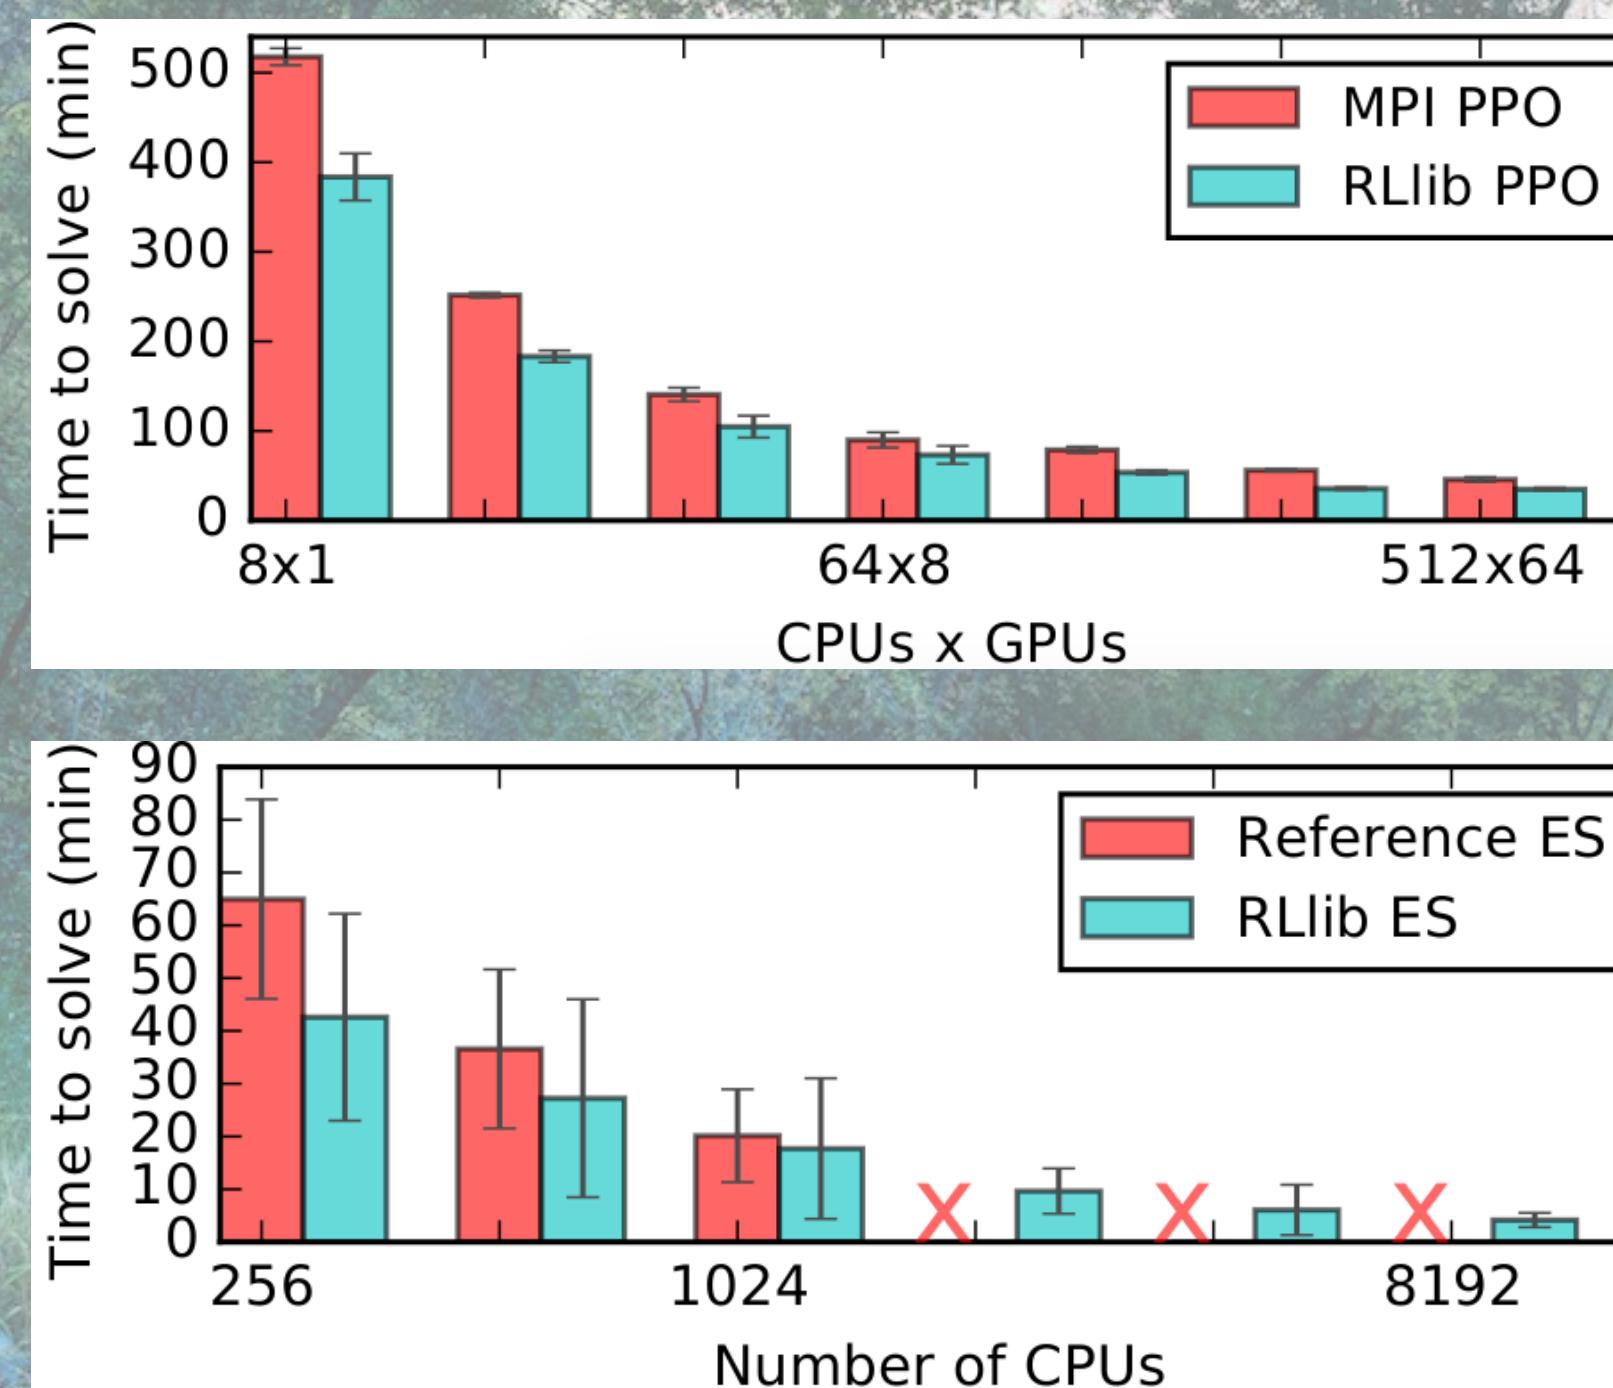


- border fight
- attack
- center ko
- nobi
- hane
- split shape

example creation:
Bob van den Hoek

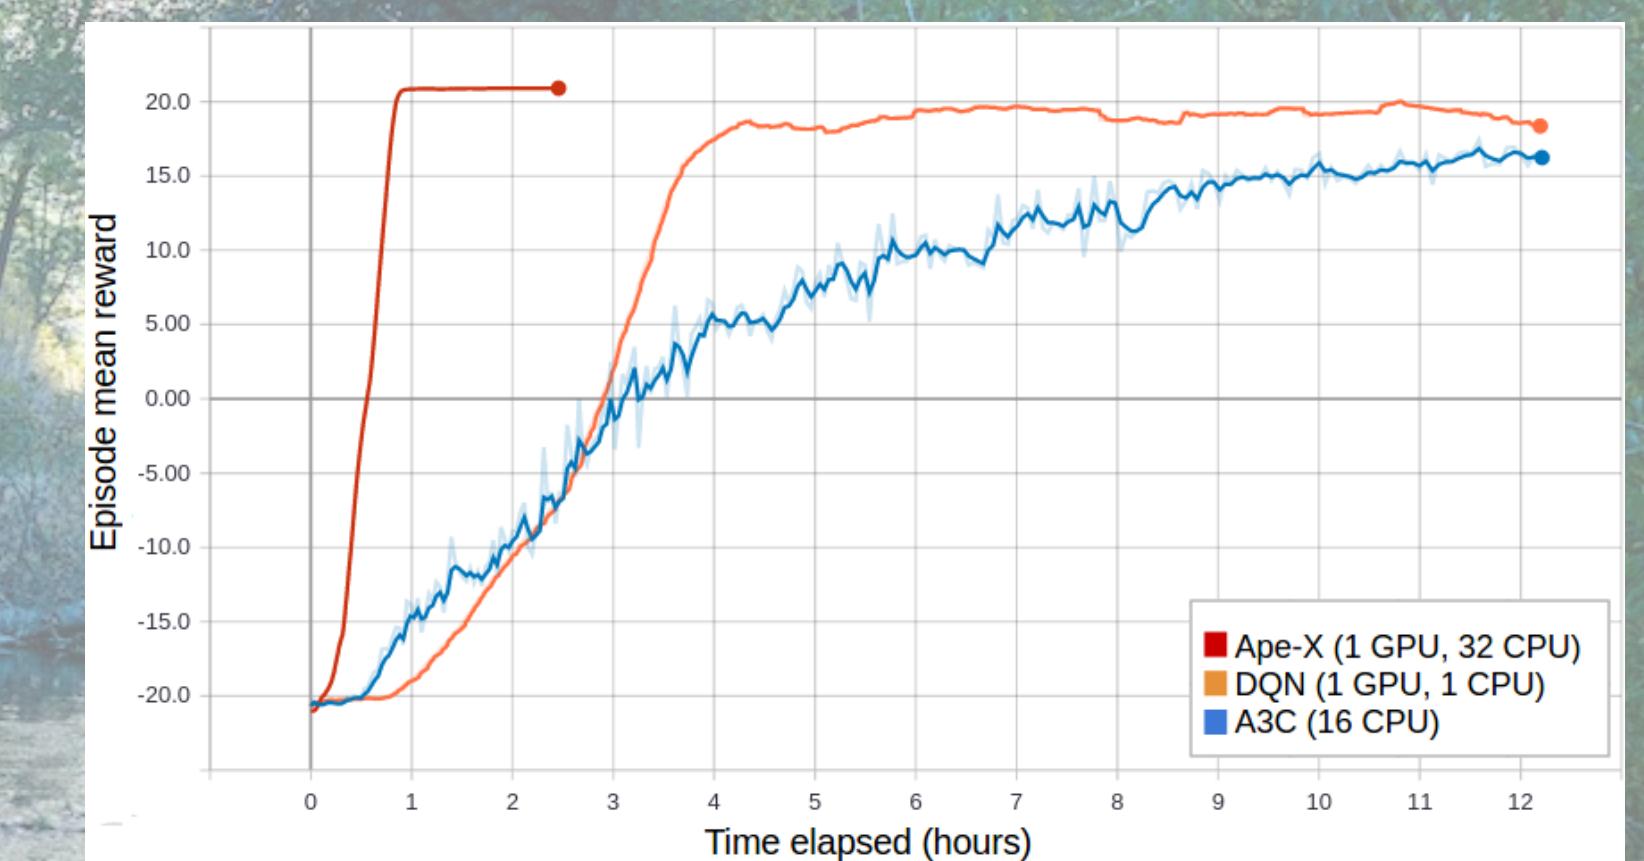
RLLib Provides a Unified Framework for Scalable RL that Doesn't Compromise on Performance

Distributed PPO



Evolution
Strategies

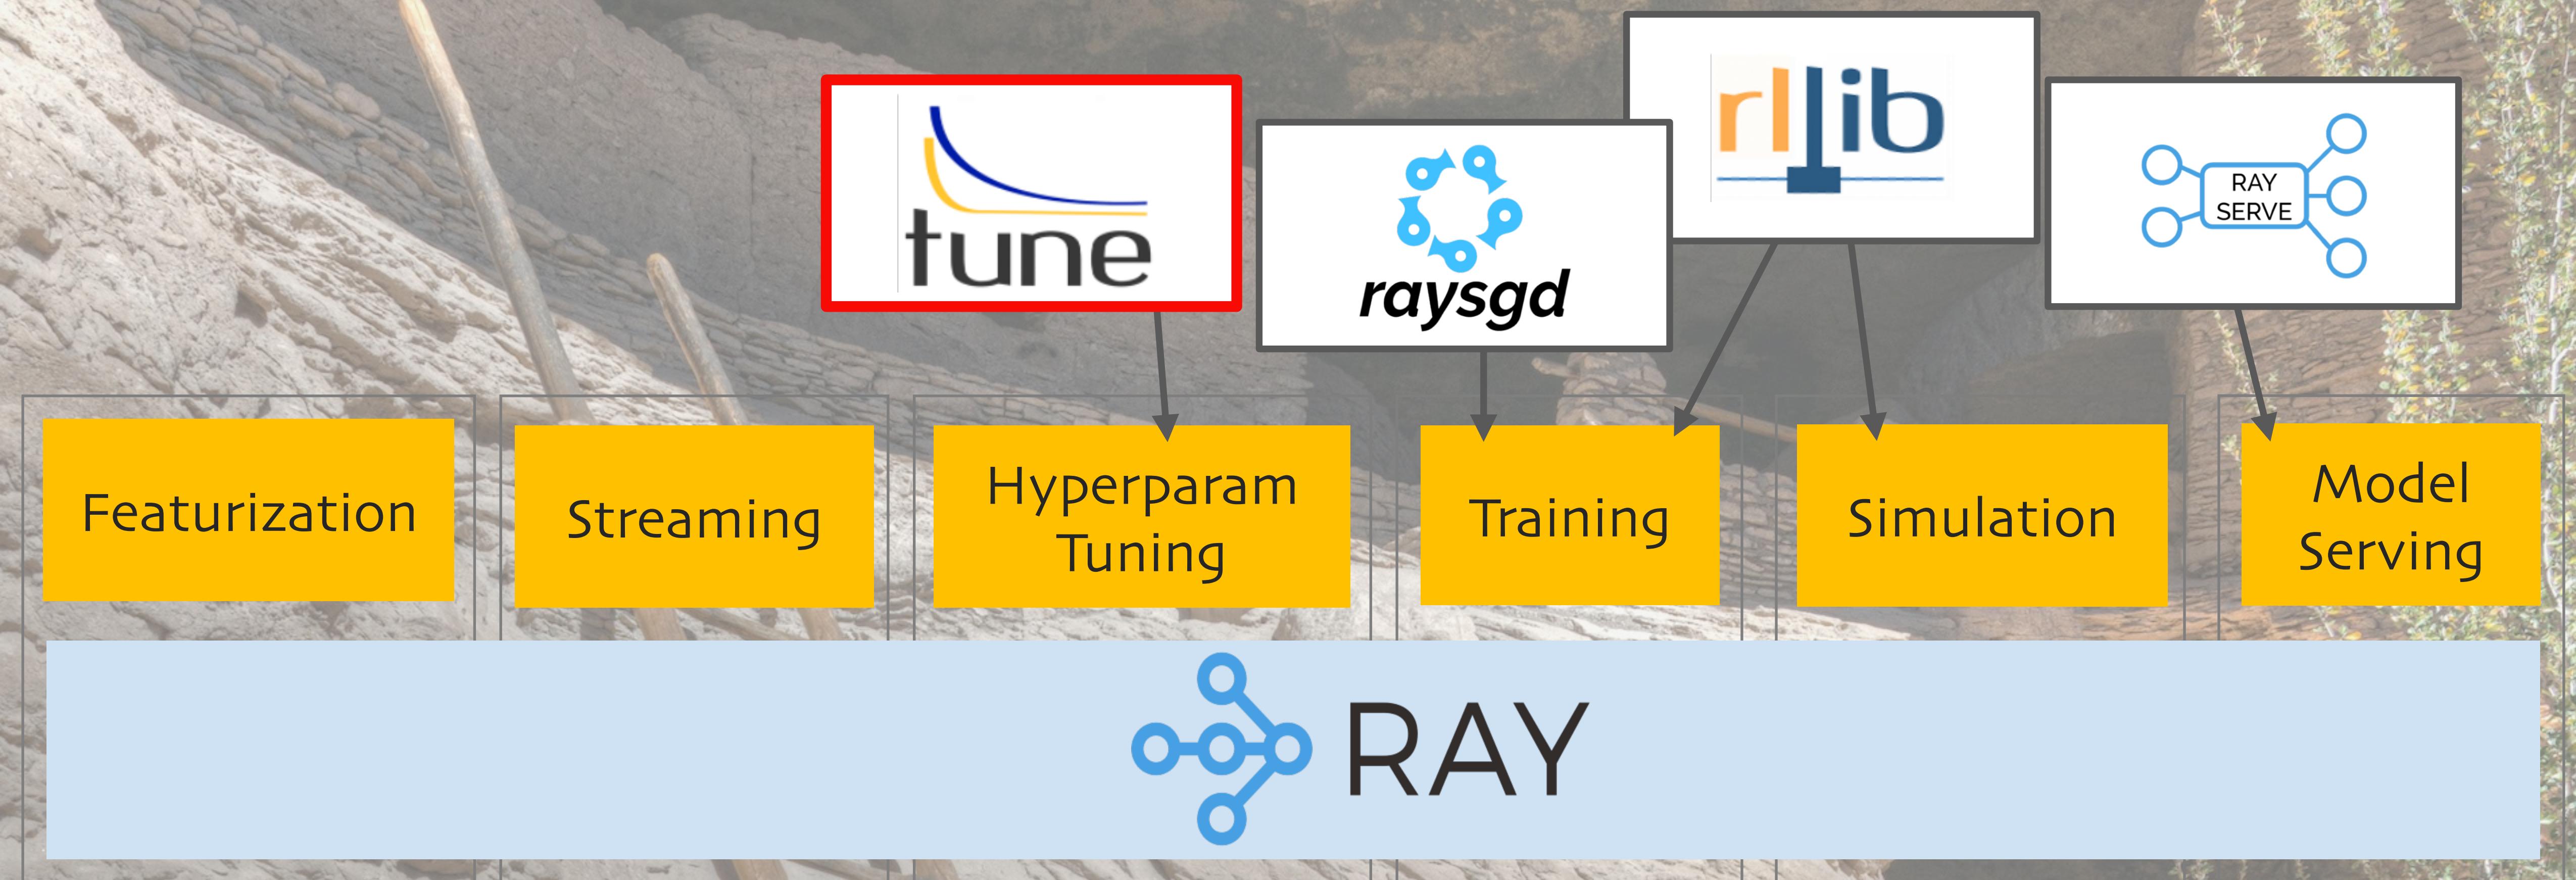
Ape-X Distributed
DQN, DDPG





@deanwampler

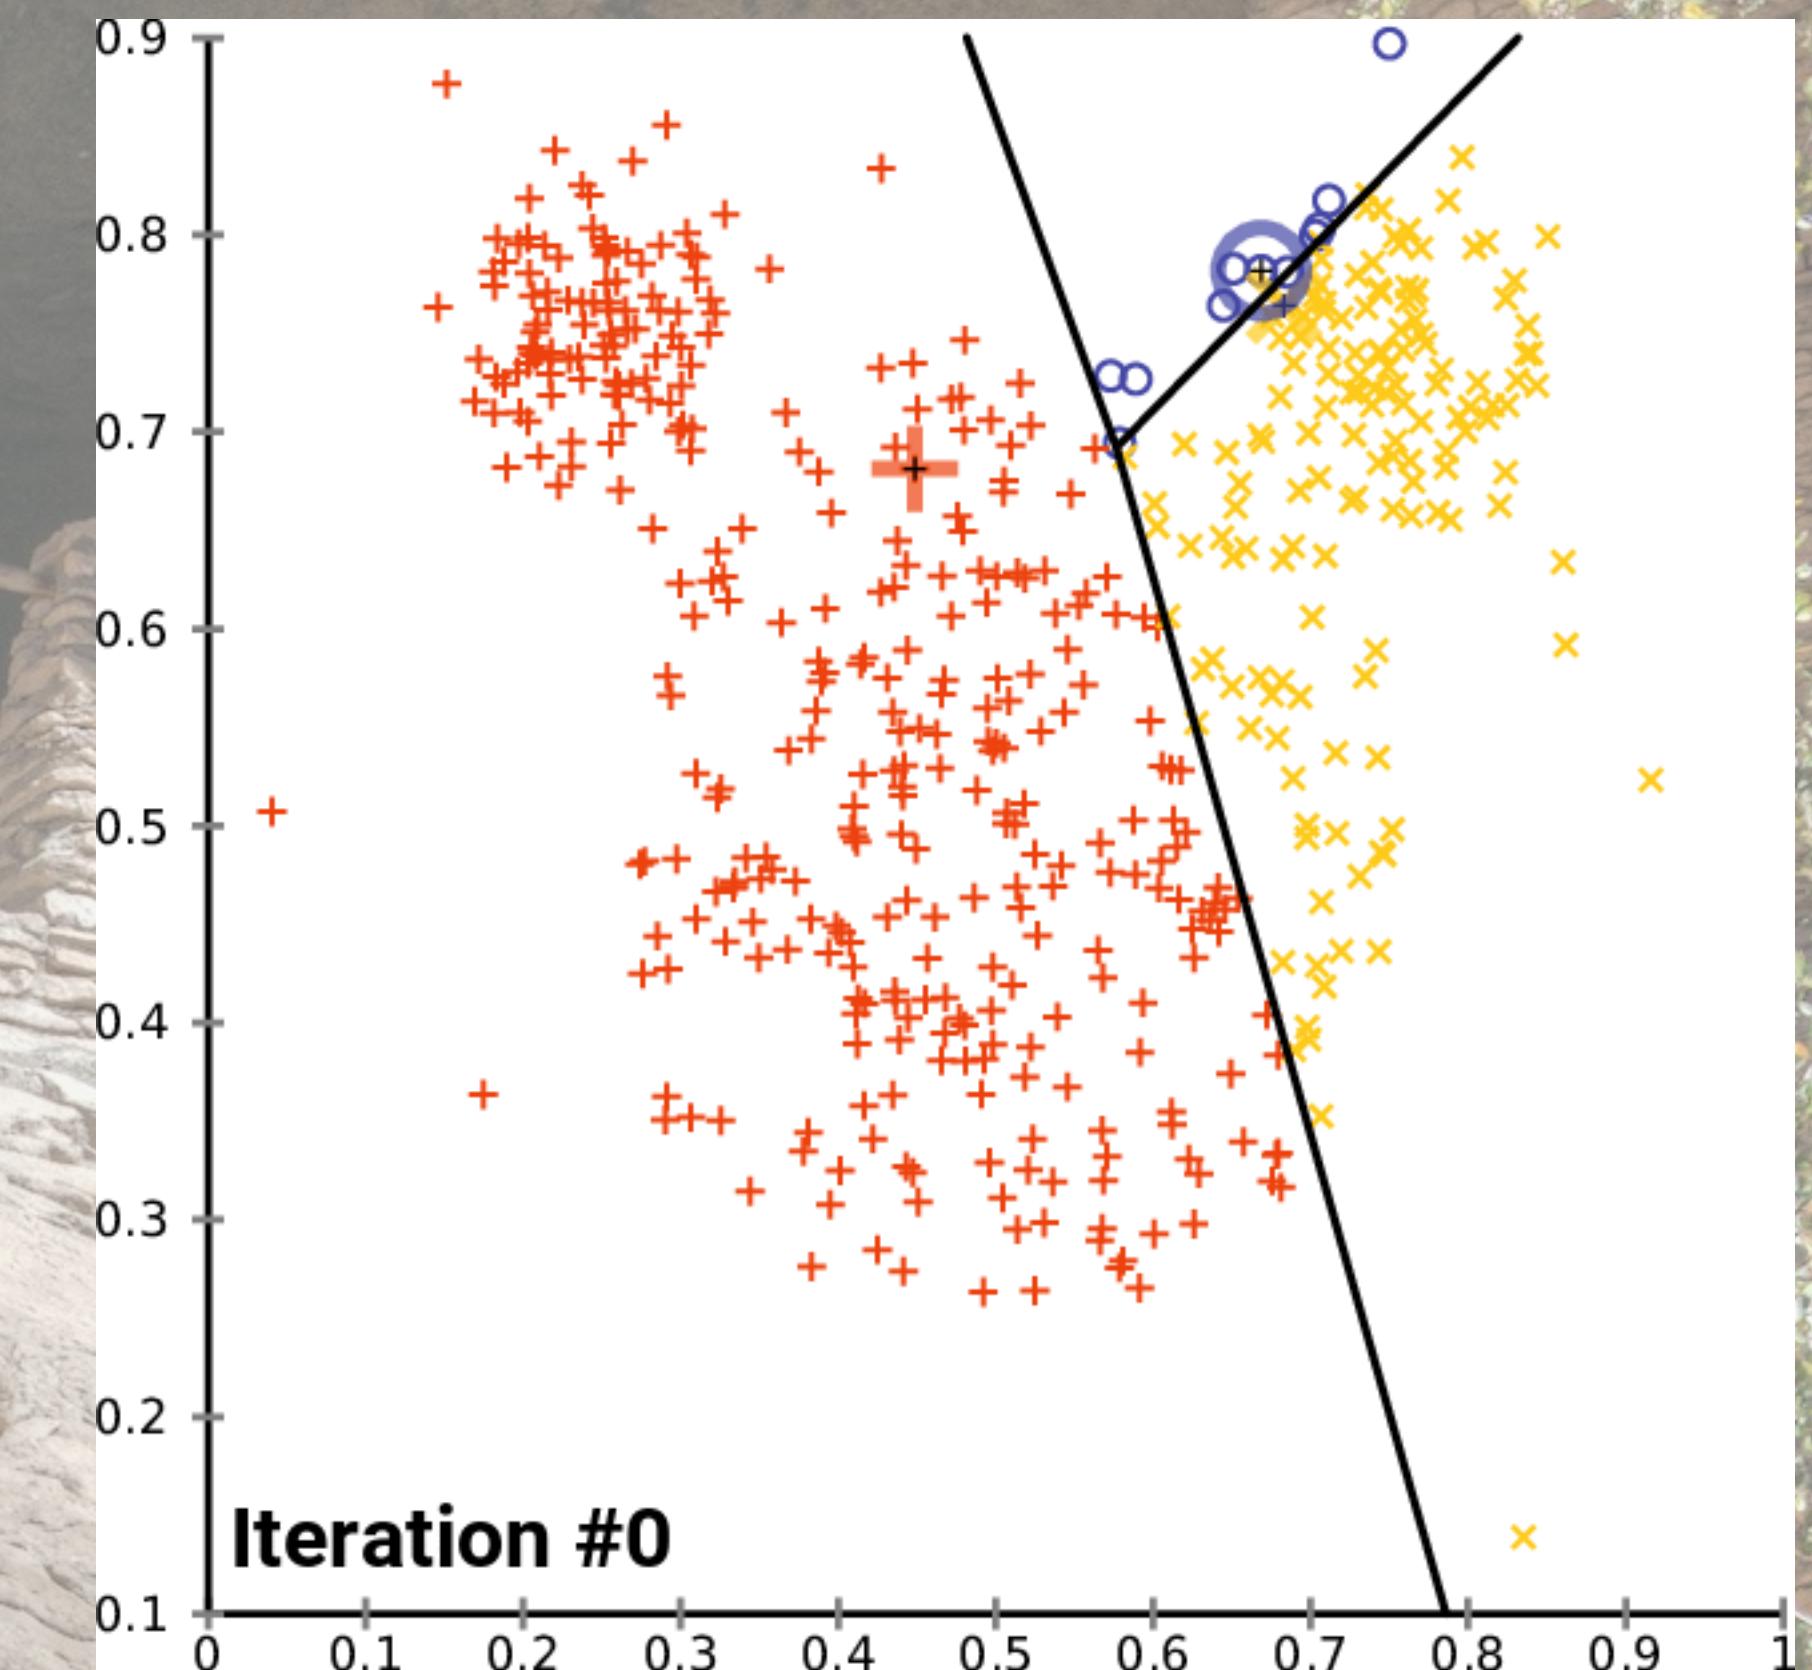
Hyperparameter Tuning - Ray Tune



What Is Hyperparameter Tuning?

Trivial example:

- What's the best value for "k" in k-means??
- k is a "hyperparameter"
- The resulting clusters are defined by "parameters"



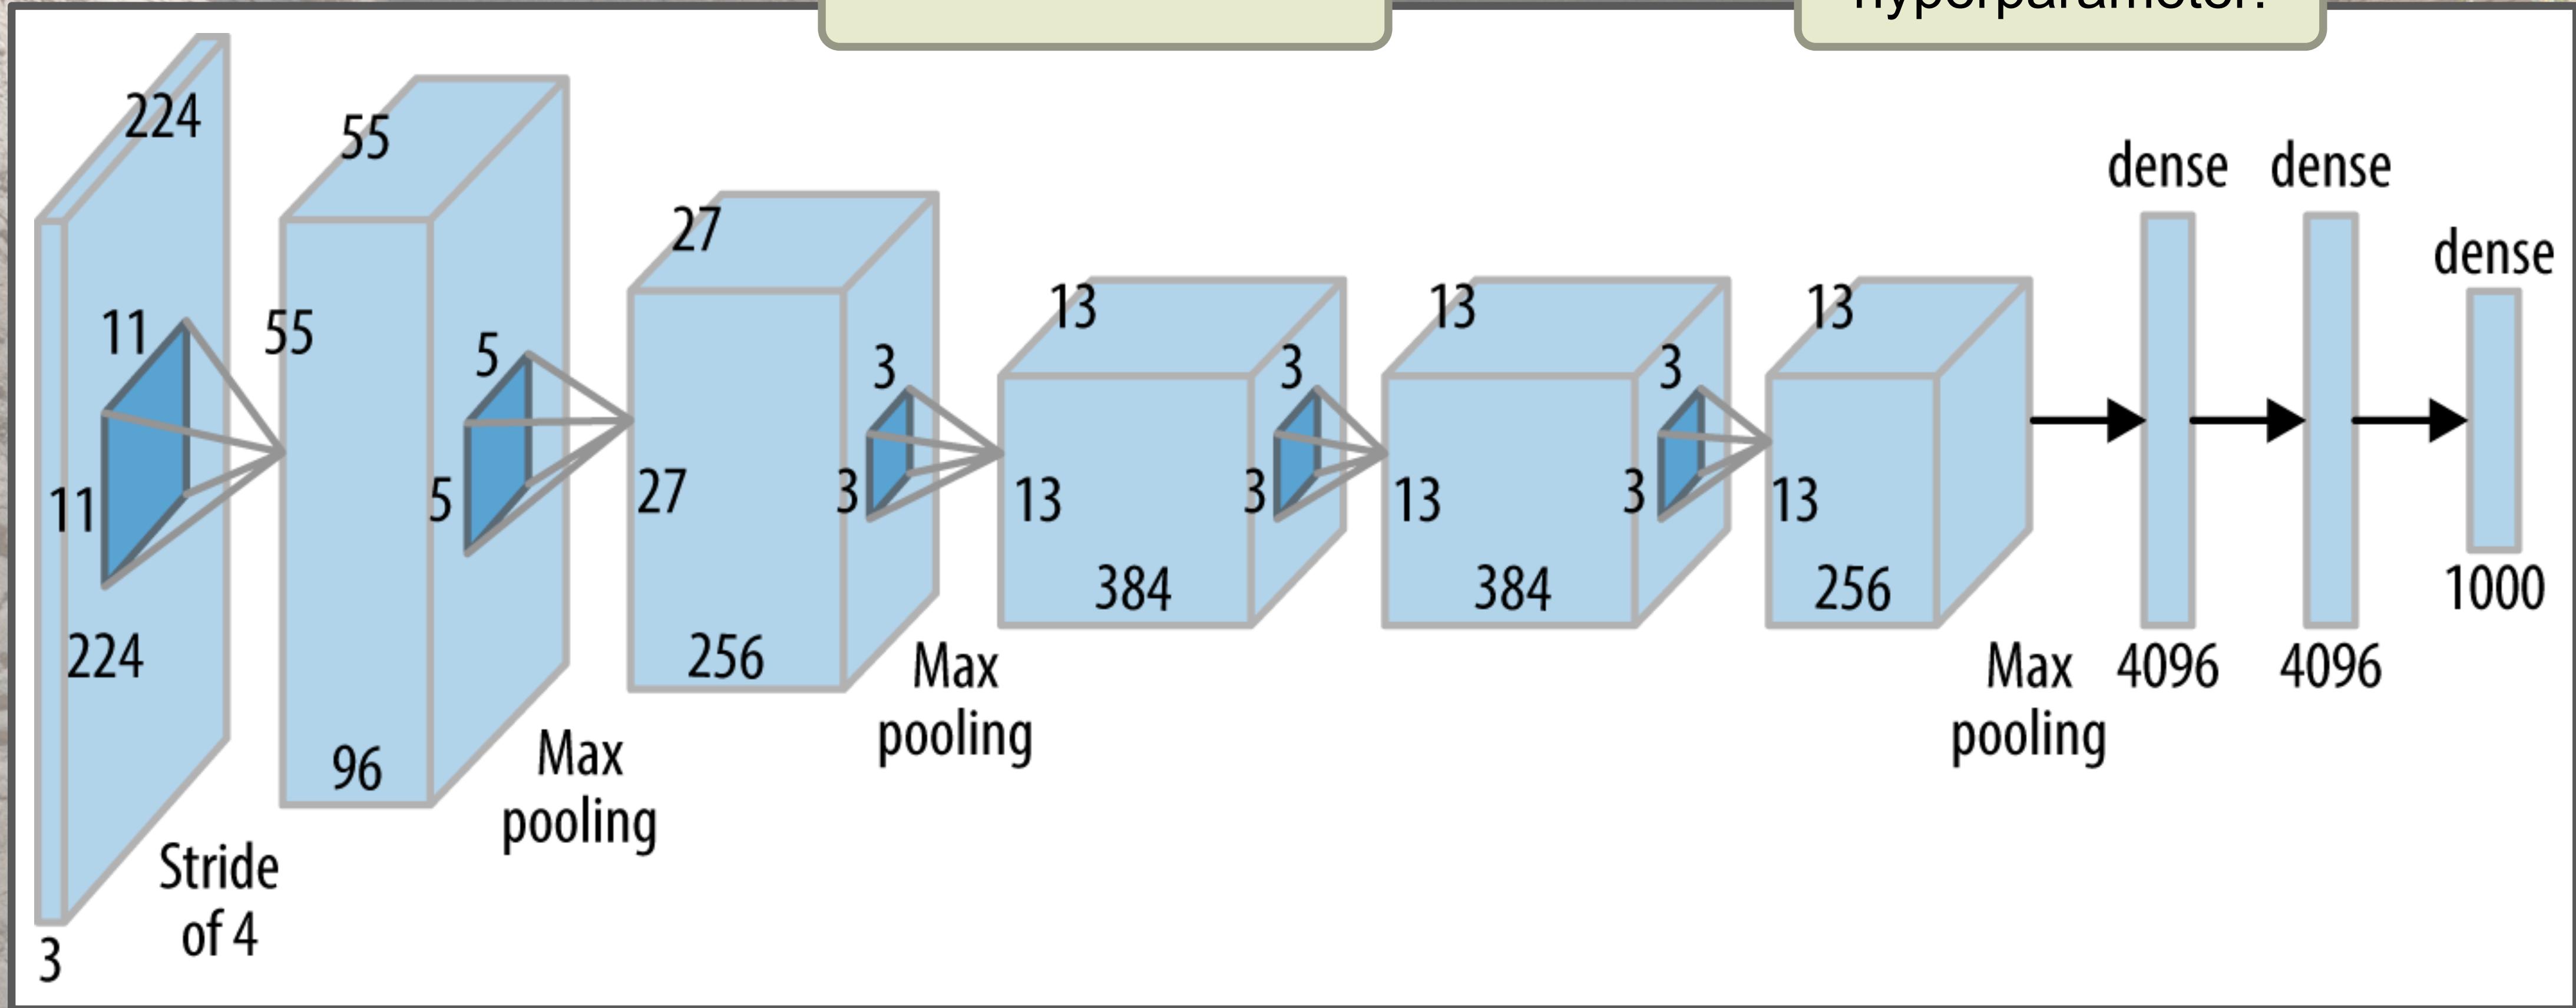
credit: https://commons.wikimedia.org/wiki/File:K-means_convergence.gif



Nontrivial Example - Neural Networks

How many layers?
What kinds of layers?

Every number
shown is a
hyperparameter!

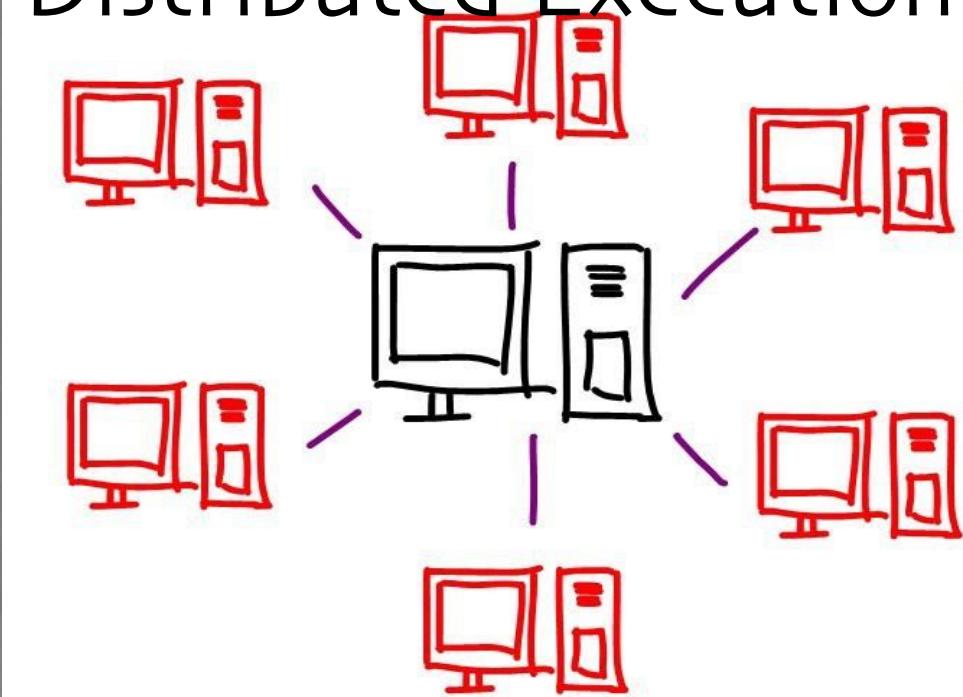


Tune is Built with Deep Learning as a Priority

Resource Aware
Scheduling



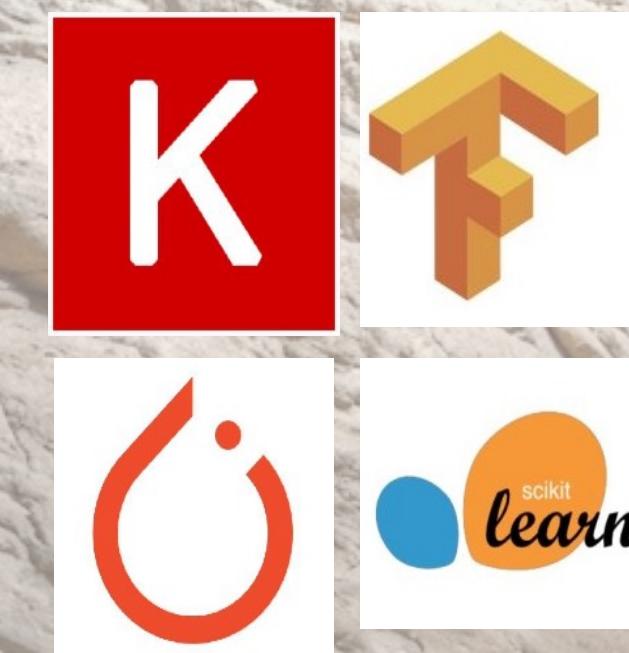
Seamless
Distributed Execution



Simple API for
new algorithms

```
class TrialScheduler:  
    def on_result(self, trial, result): ...  
    def choose_trial_to_run(self): ...
```

Framework Agnostic

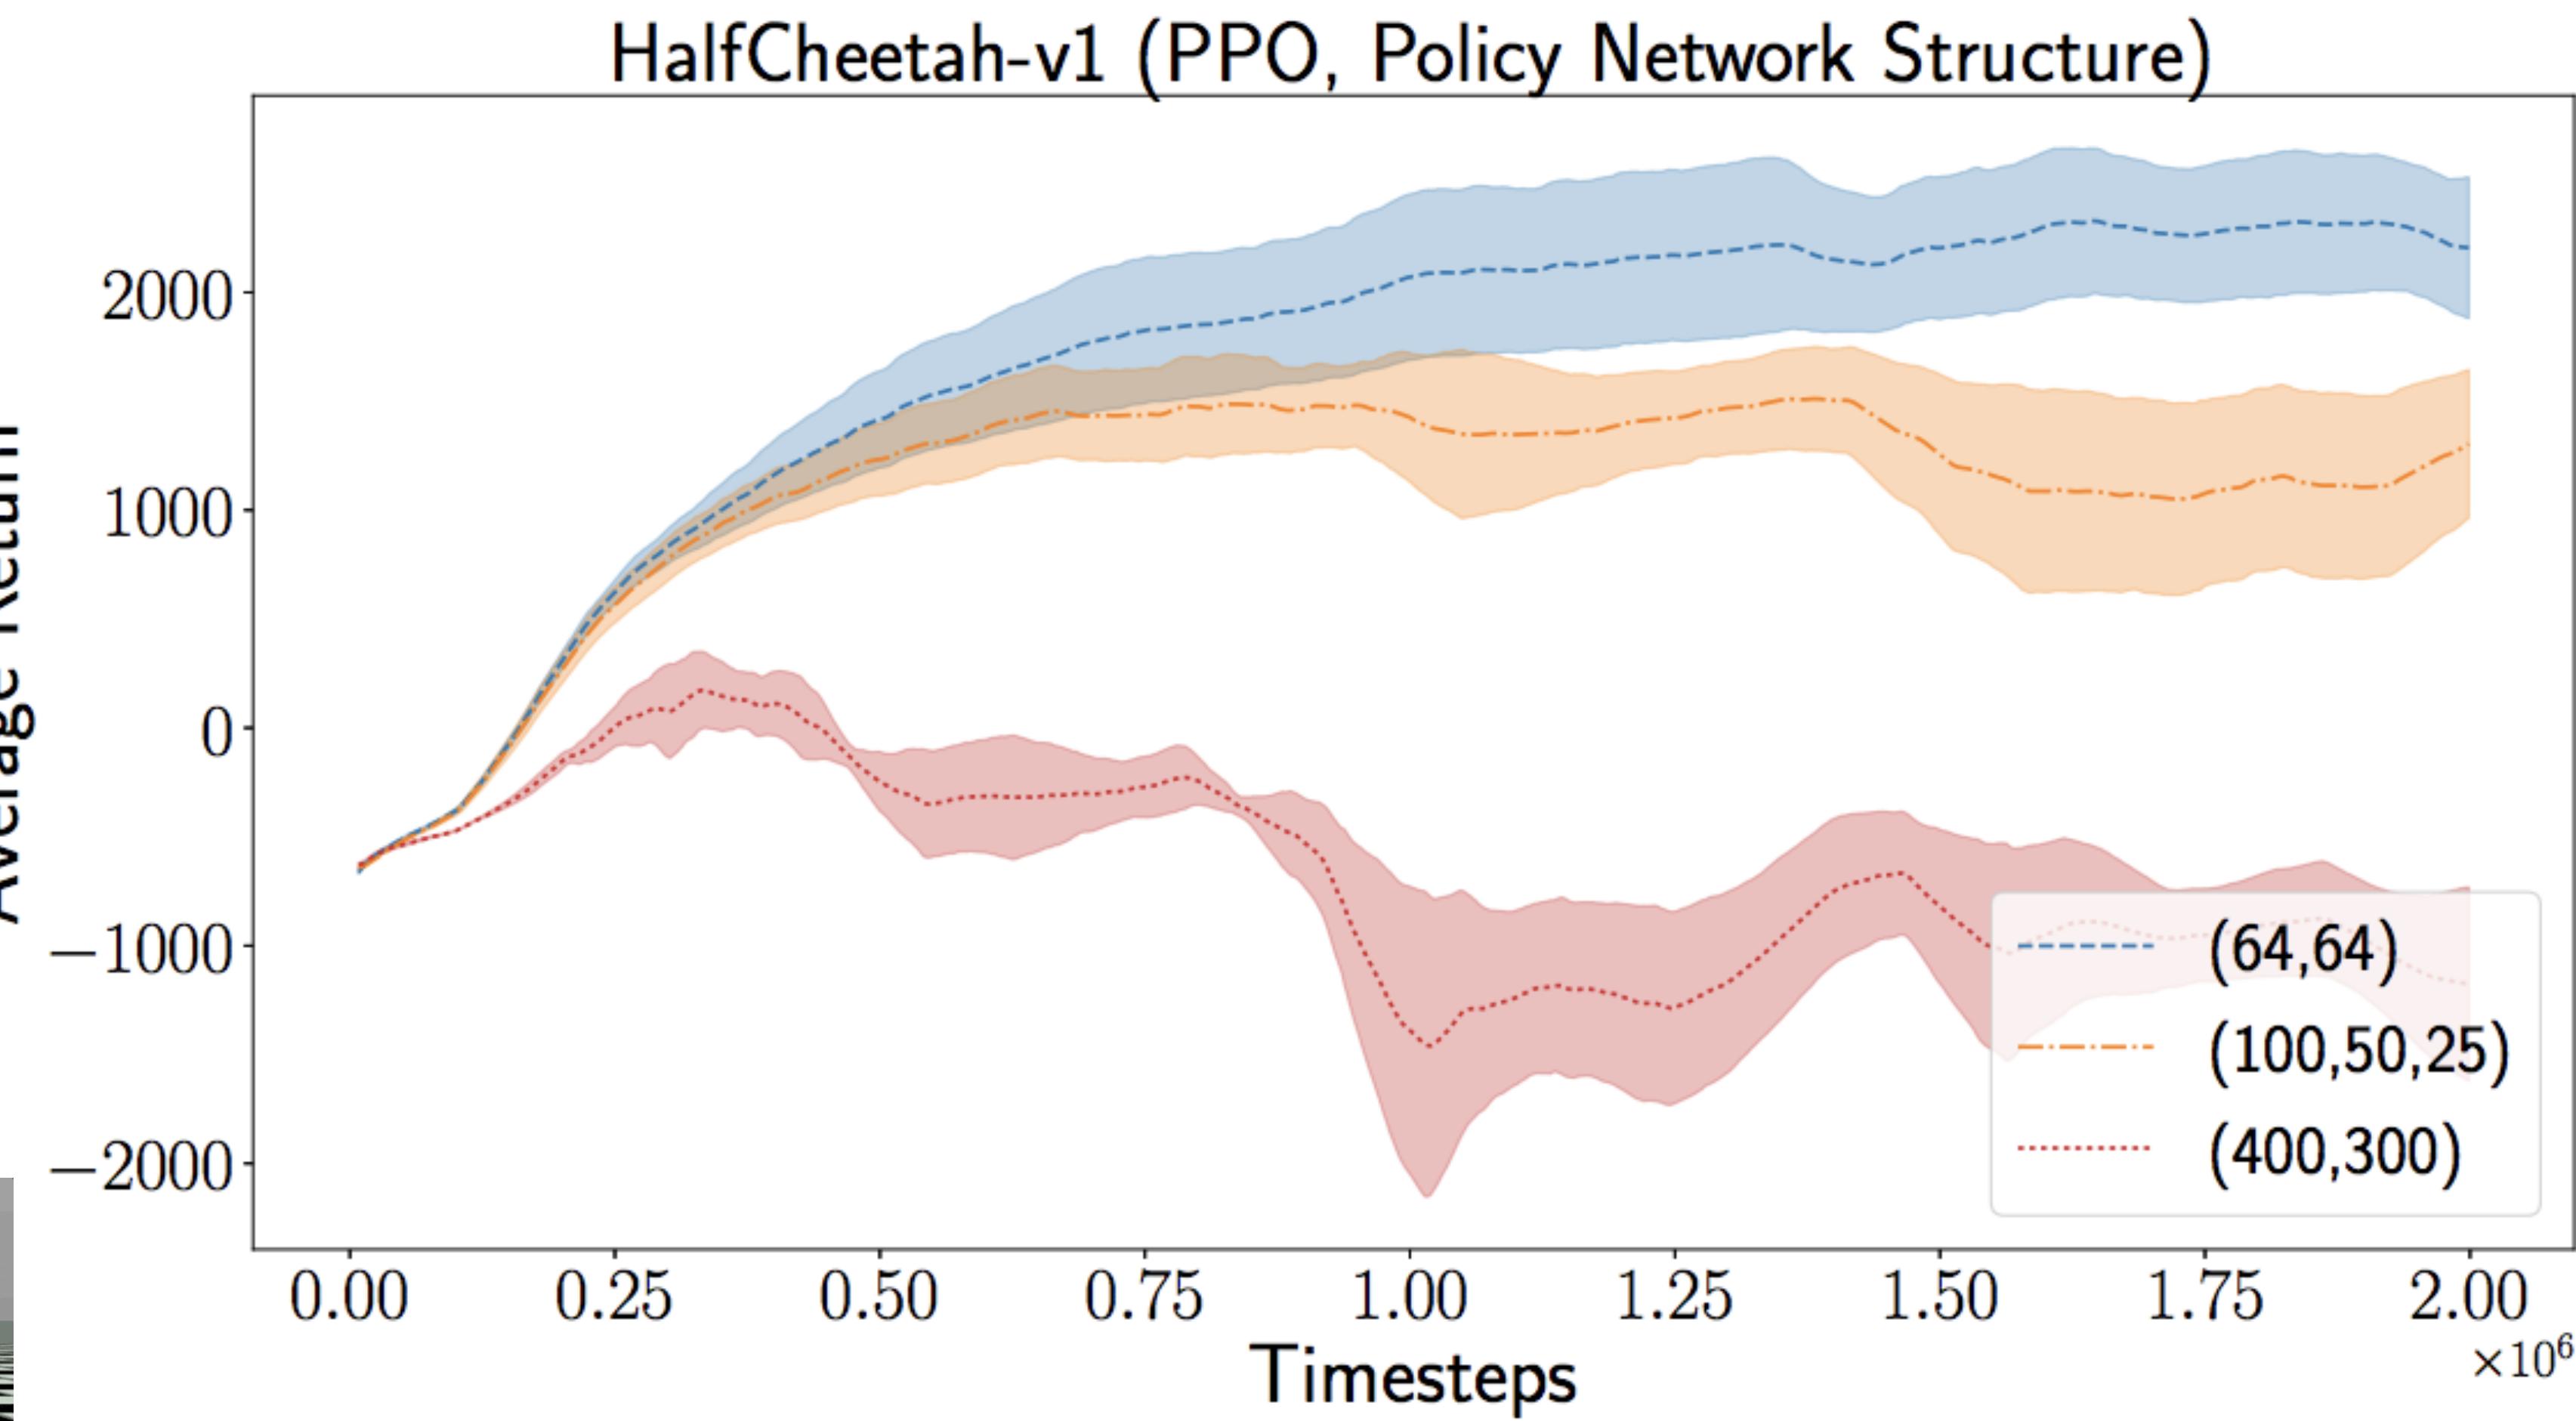


tune.io

@deanwampler



Hyperparameters Are Important for Performance

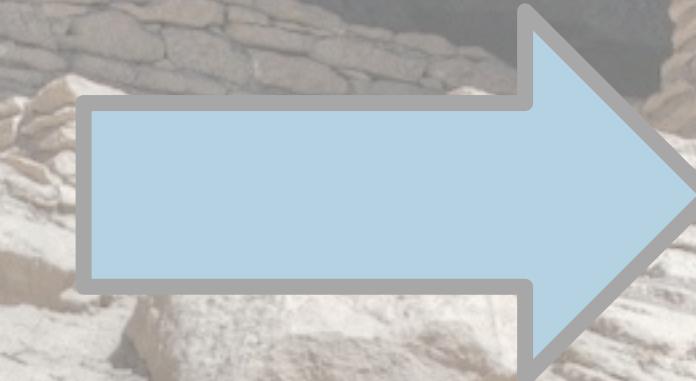


Why We Need a Framework for Tuning Hyperparameters

We want the best model

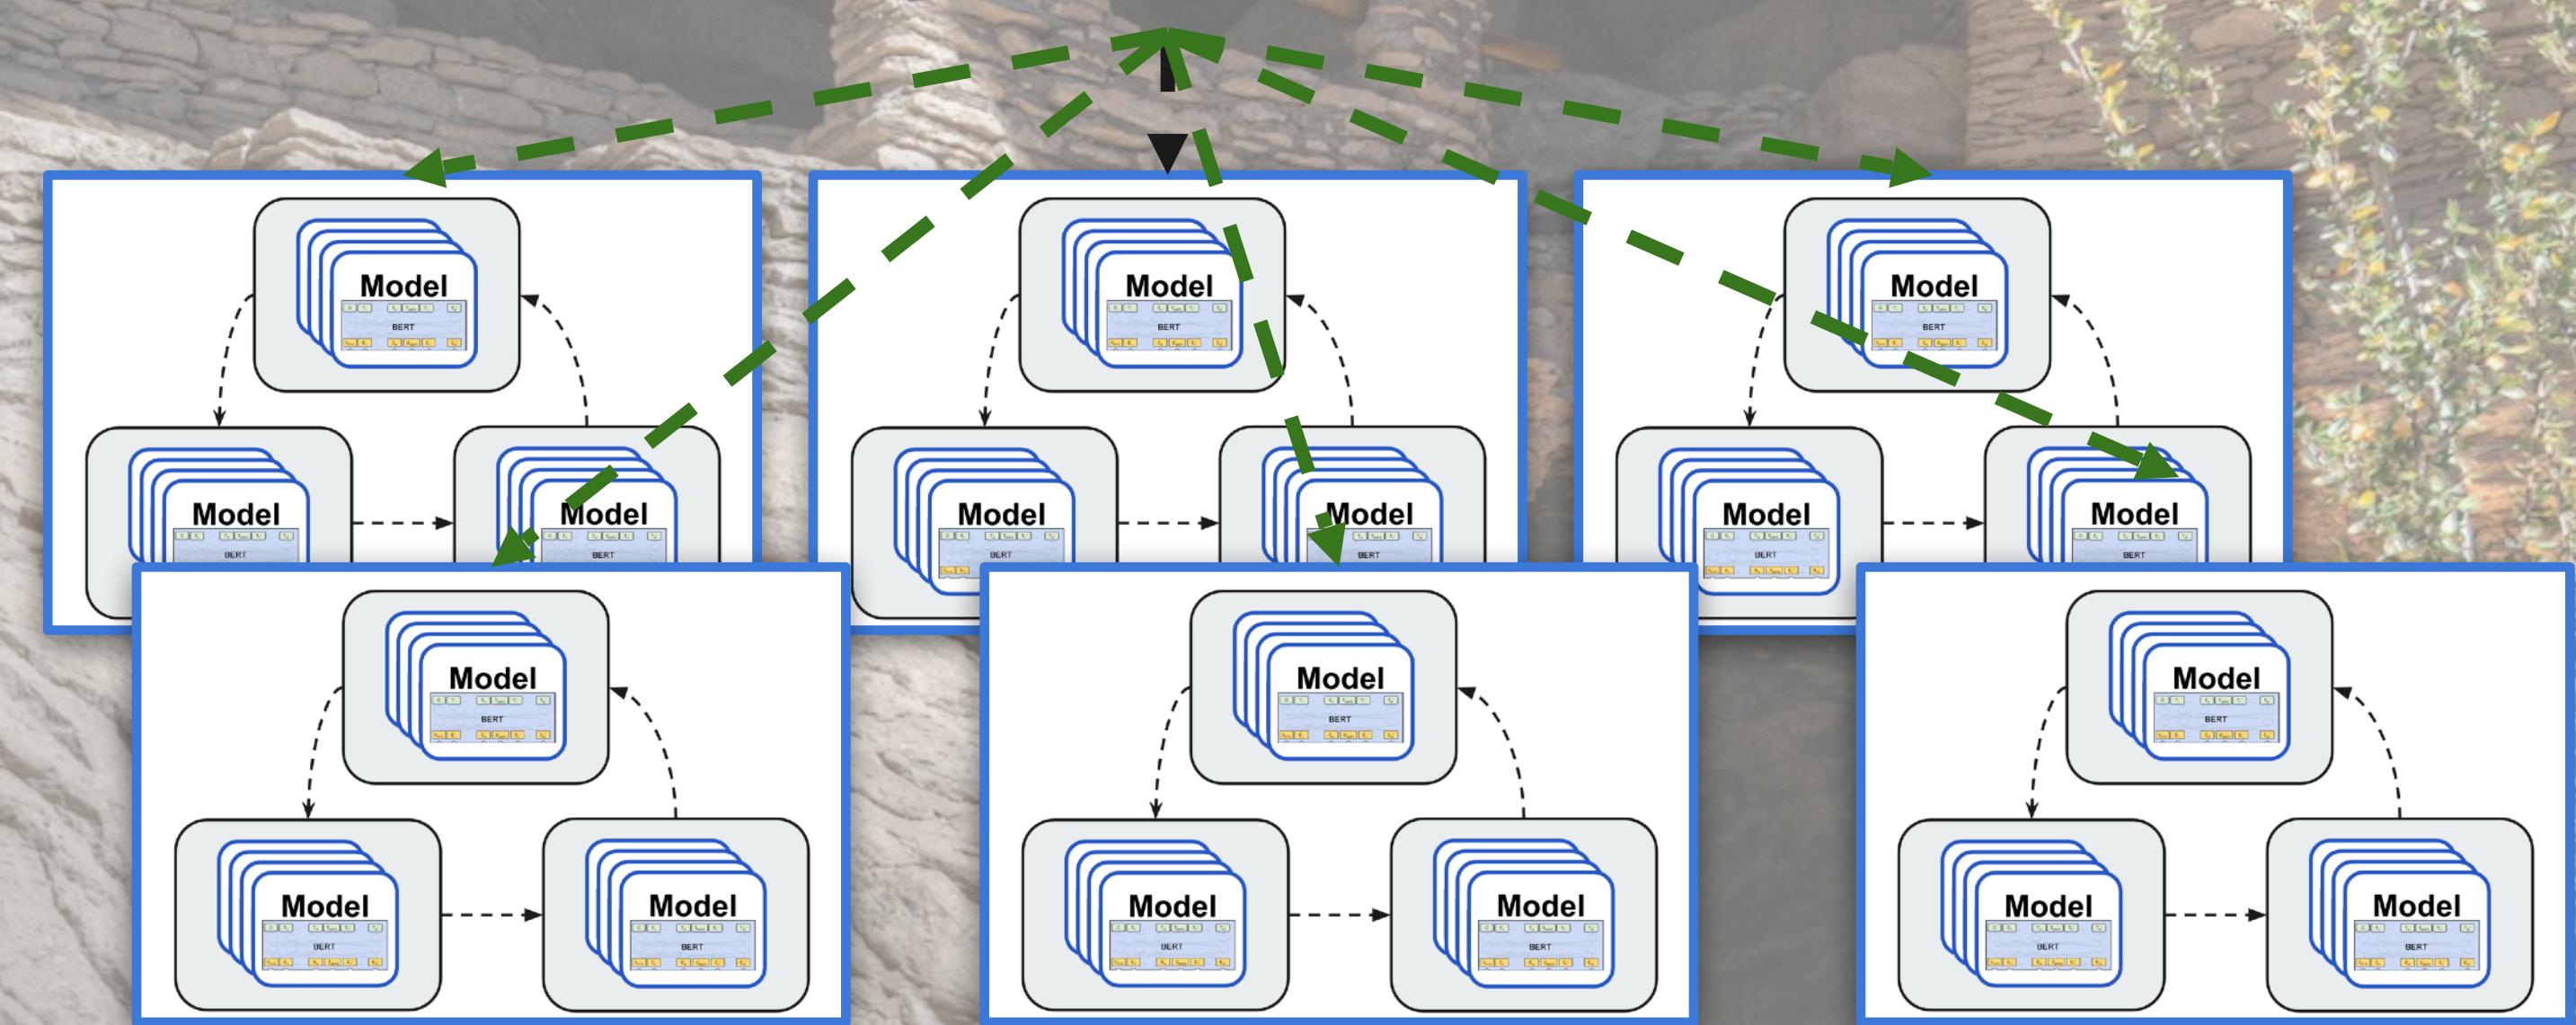
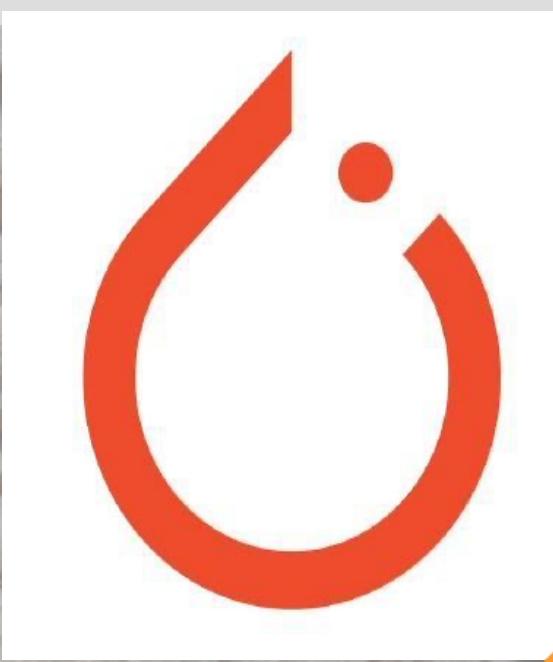
Resources are expensive

Model training is time-consuming

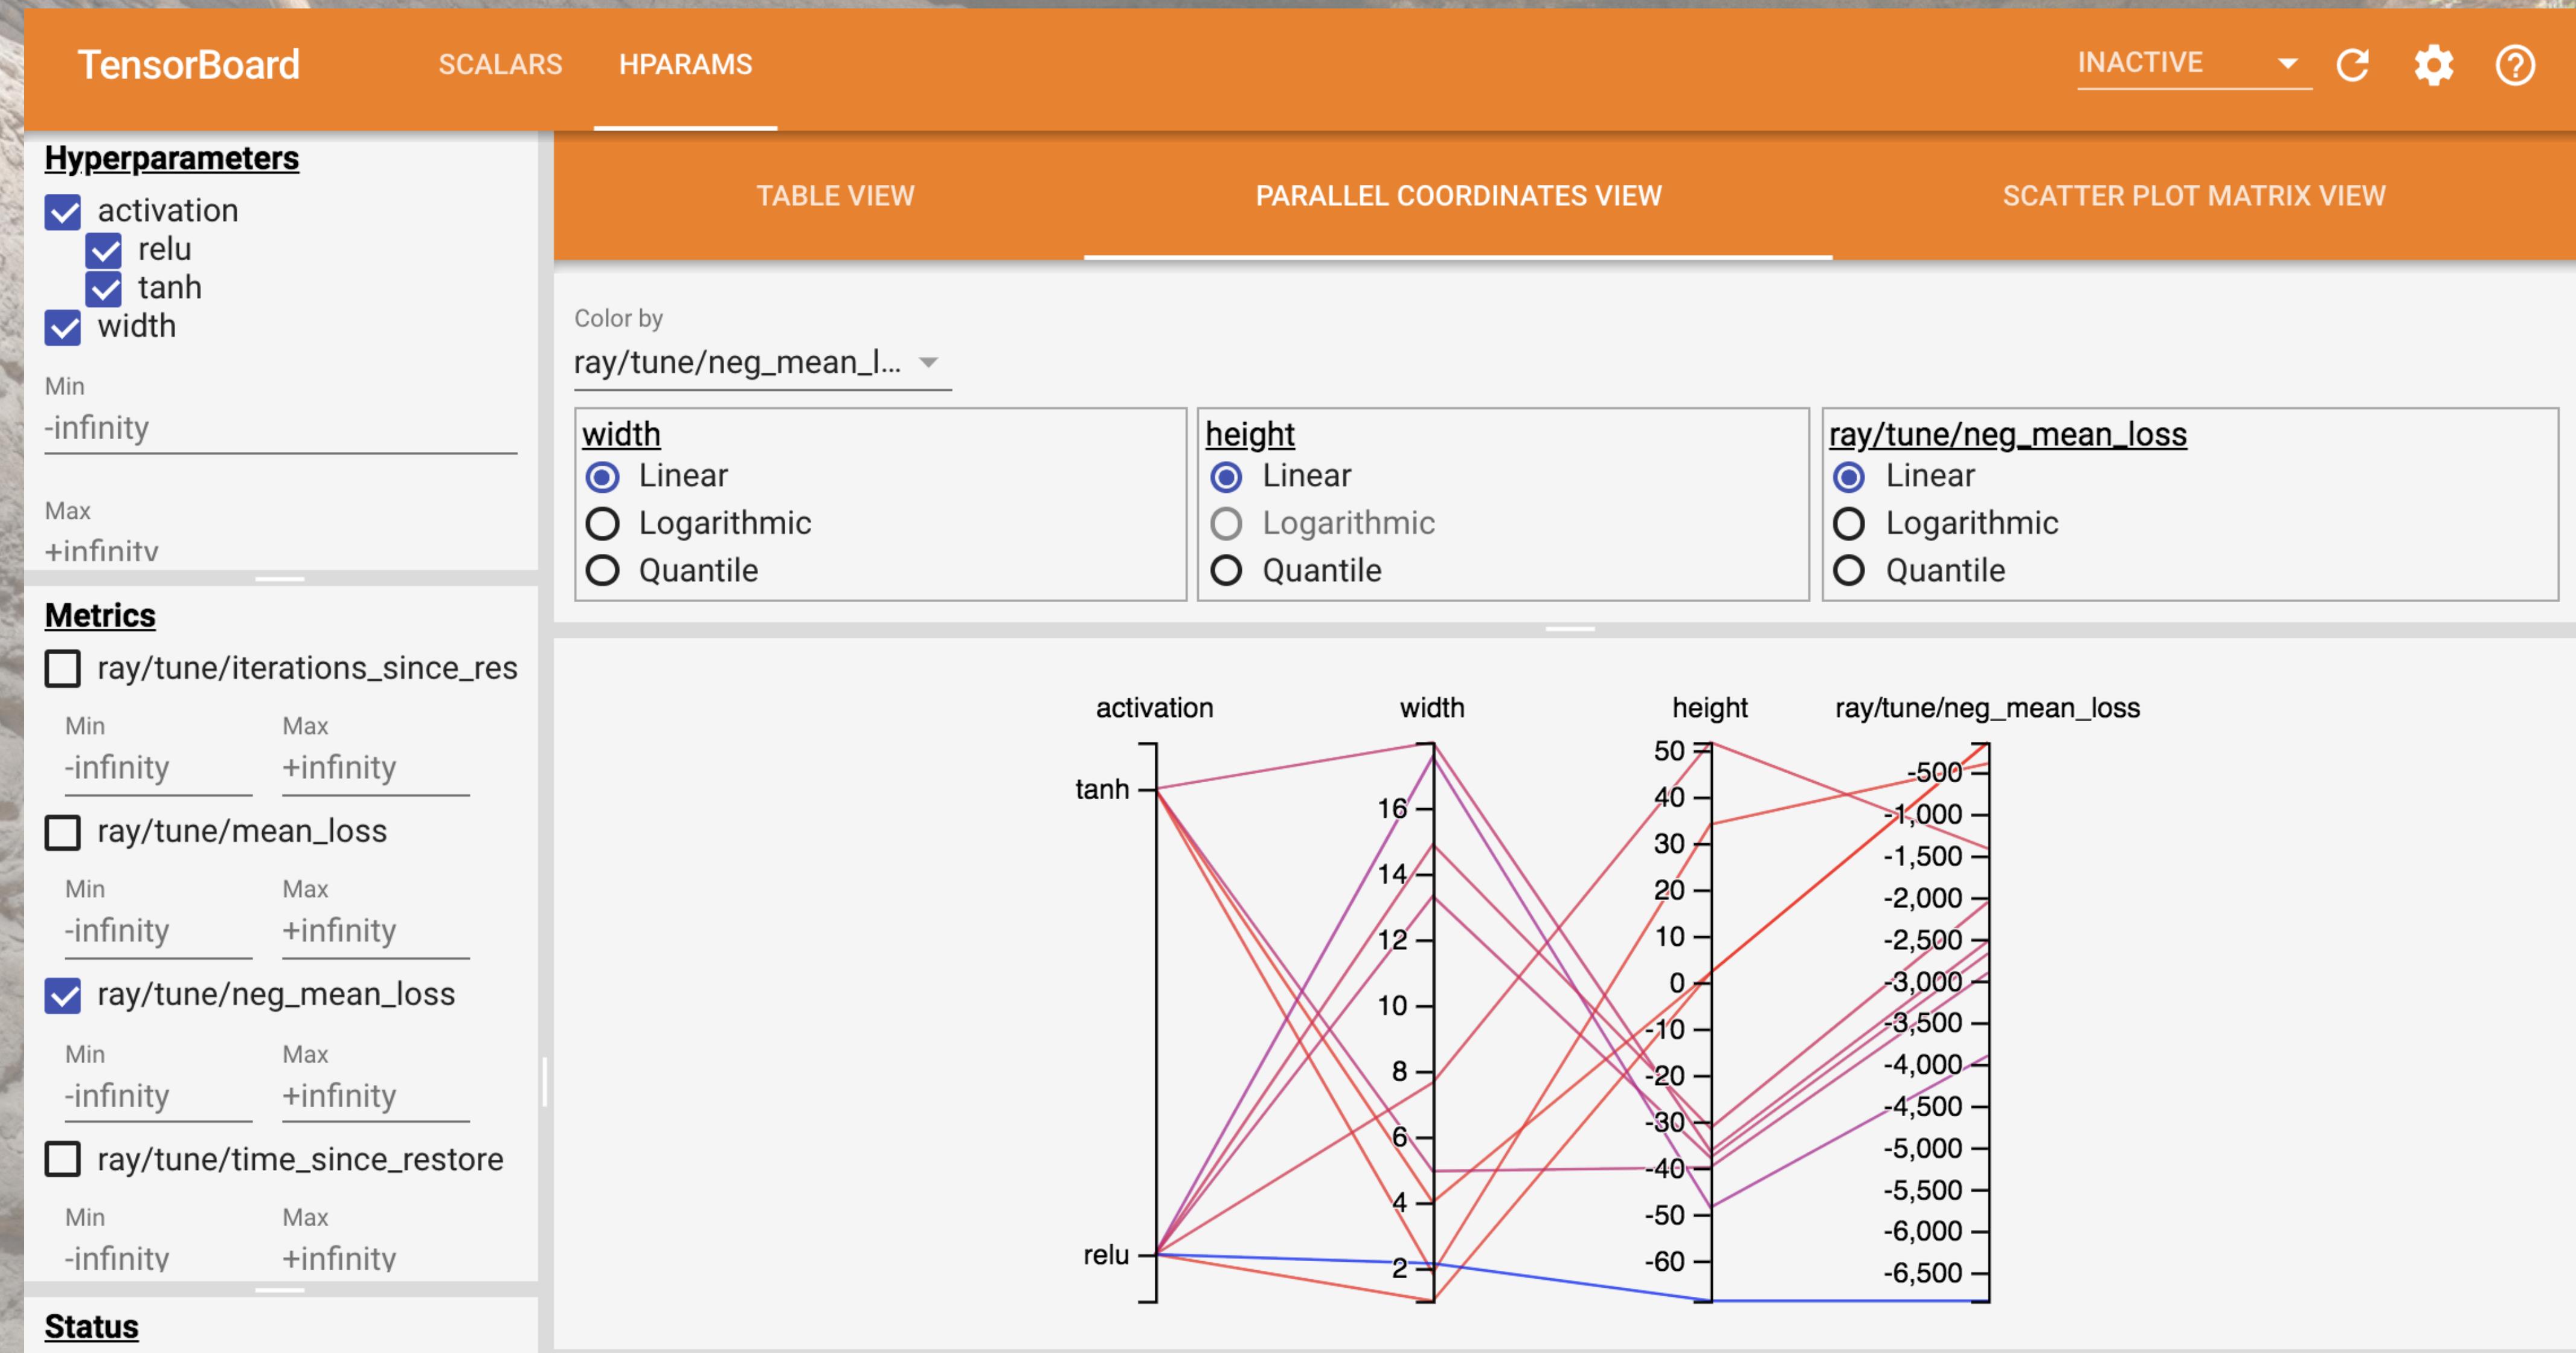


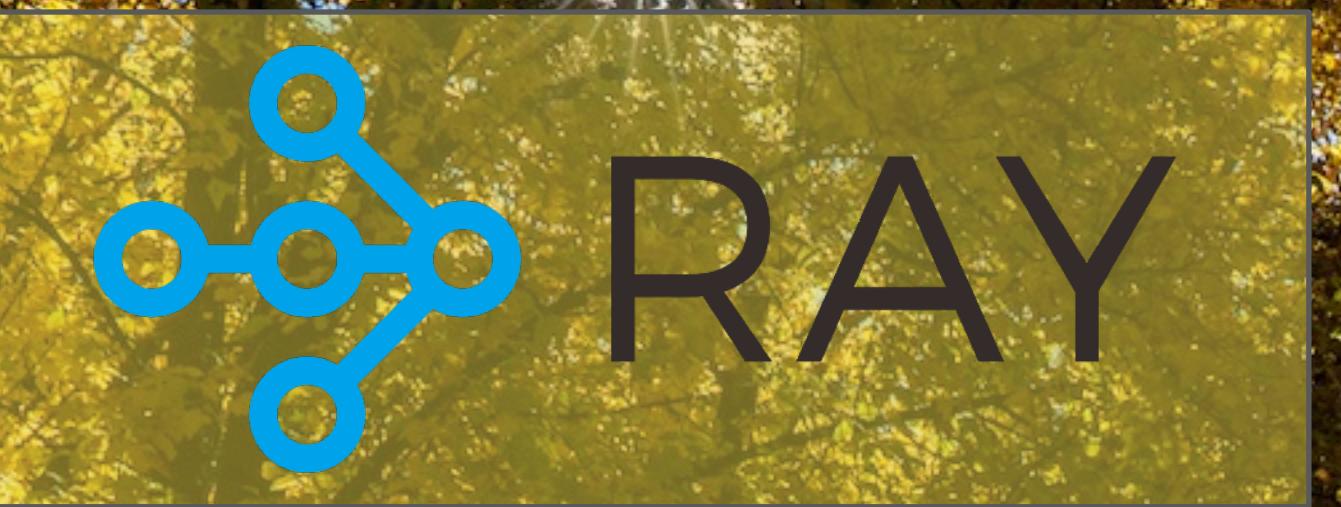
Tuning + Distributed Training

```
tune.run(PytorchTrainable,  
 config={  
     "model_creator": PretrainBERT,  
     "data_creator": create_data_loader,  
     "use_gpu": True,  
     "num_replicas": 8,  
     "lr": tune.uniform(0.001, 0.1)  
 },  
 num_samples=100,  
 search_alg=BayesianOptimization()
```



Native Integration with TensorBoard HParams





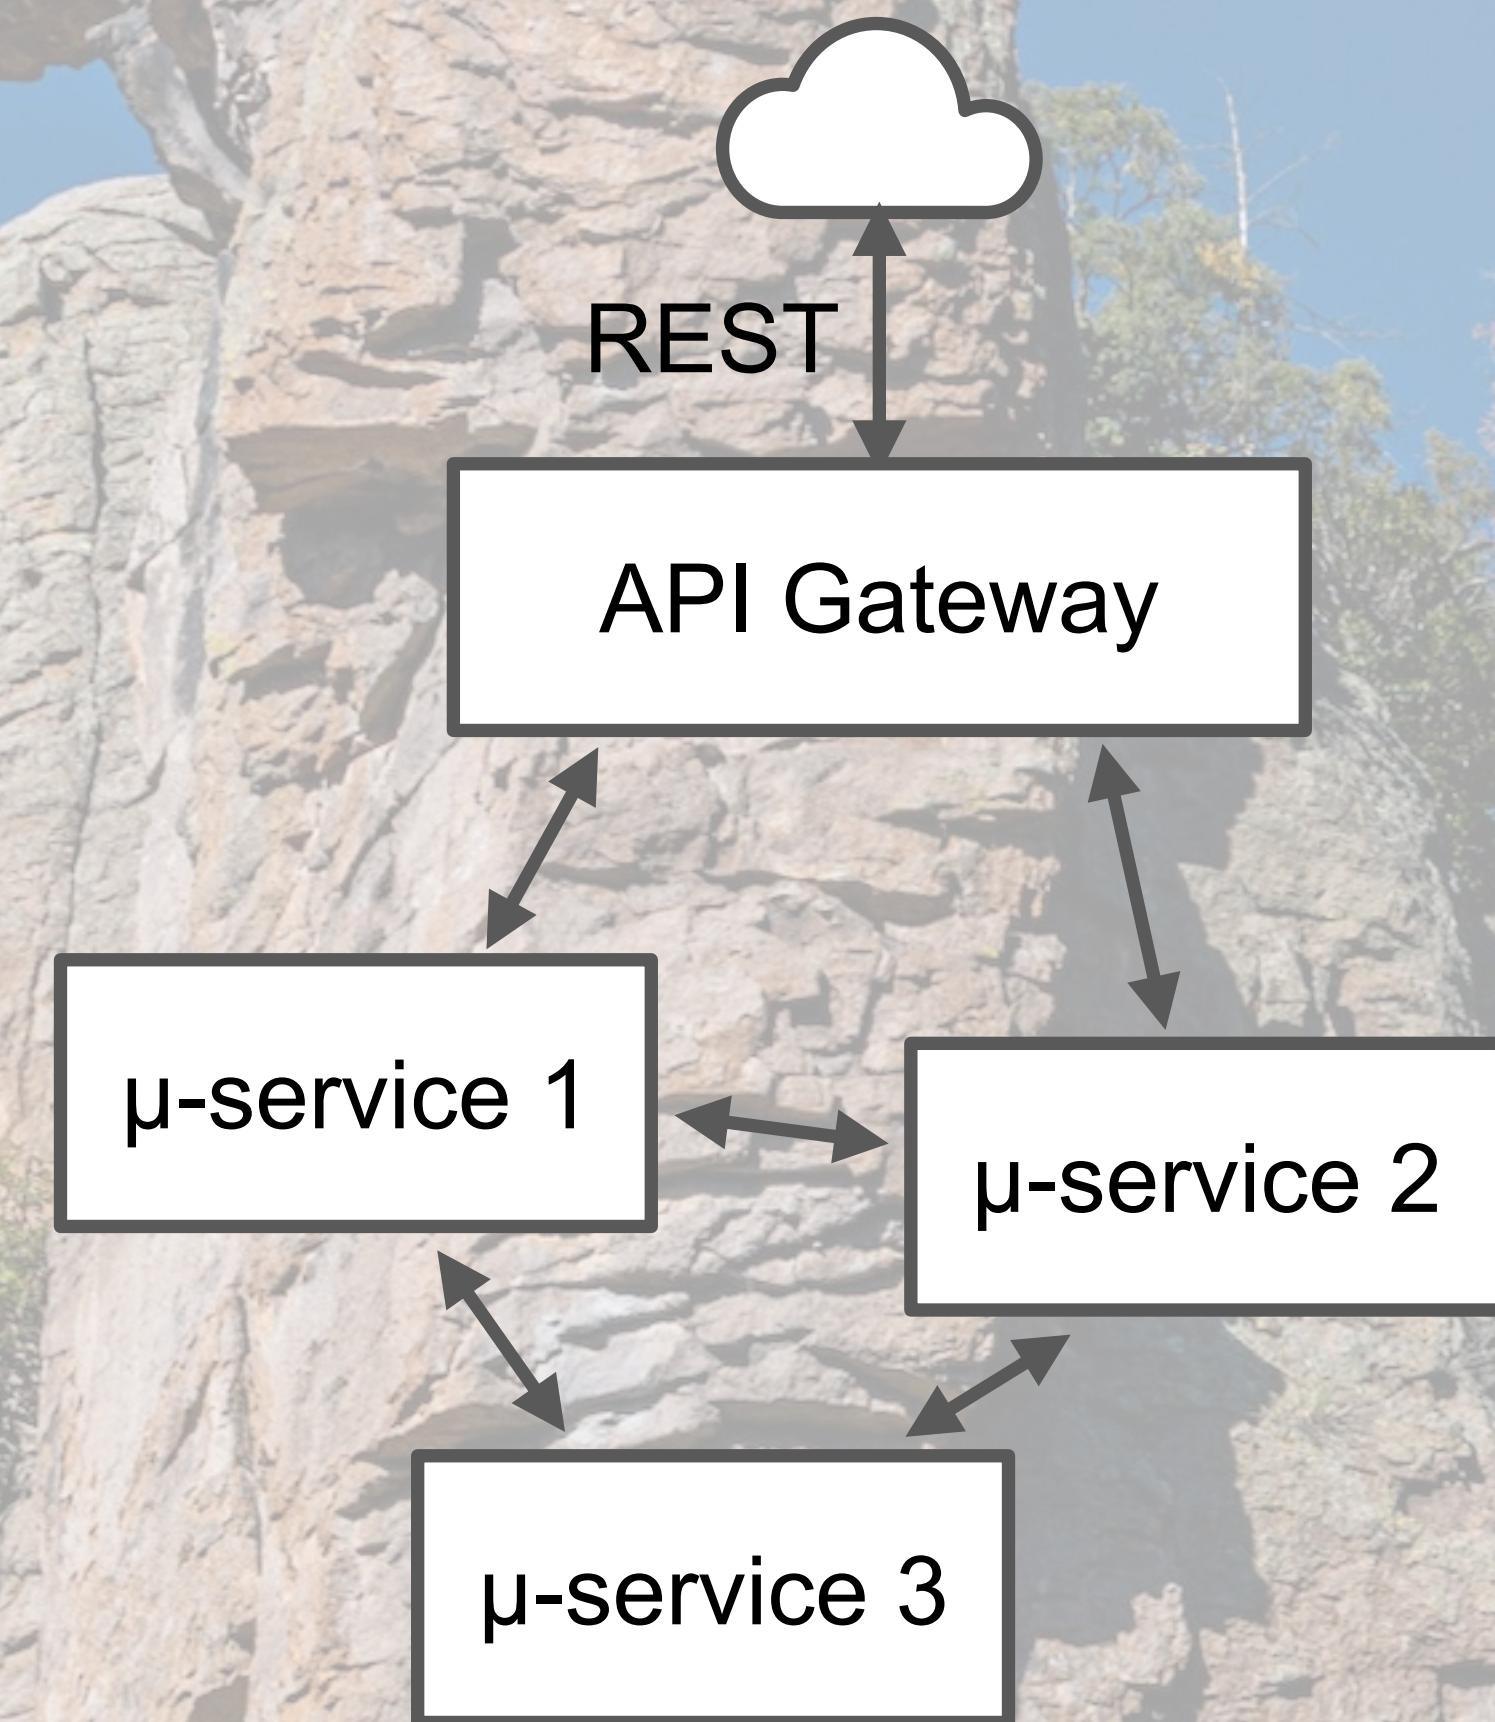
What about Ray
for Microservices?



@deanwampler

What Are Microservices?

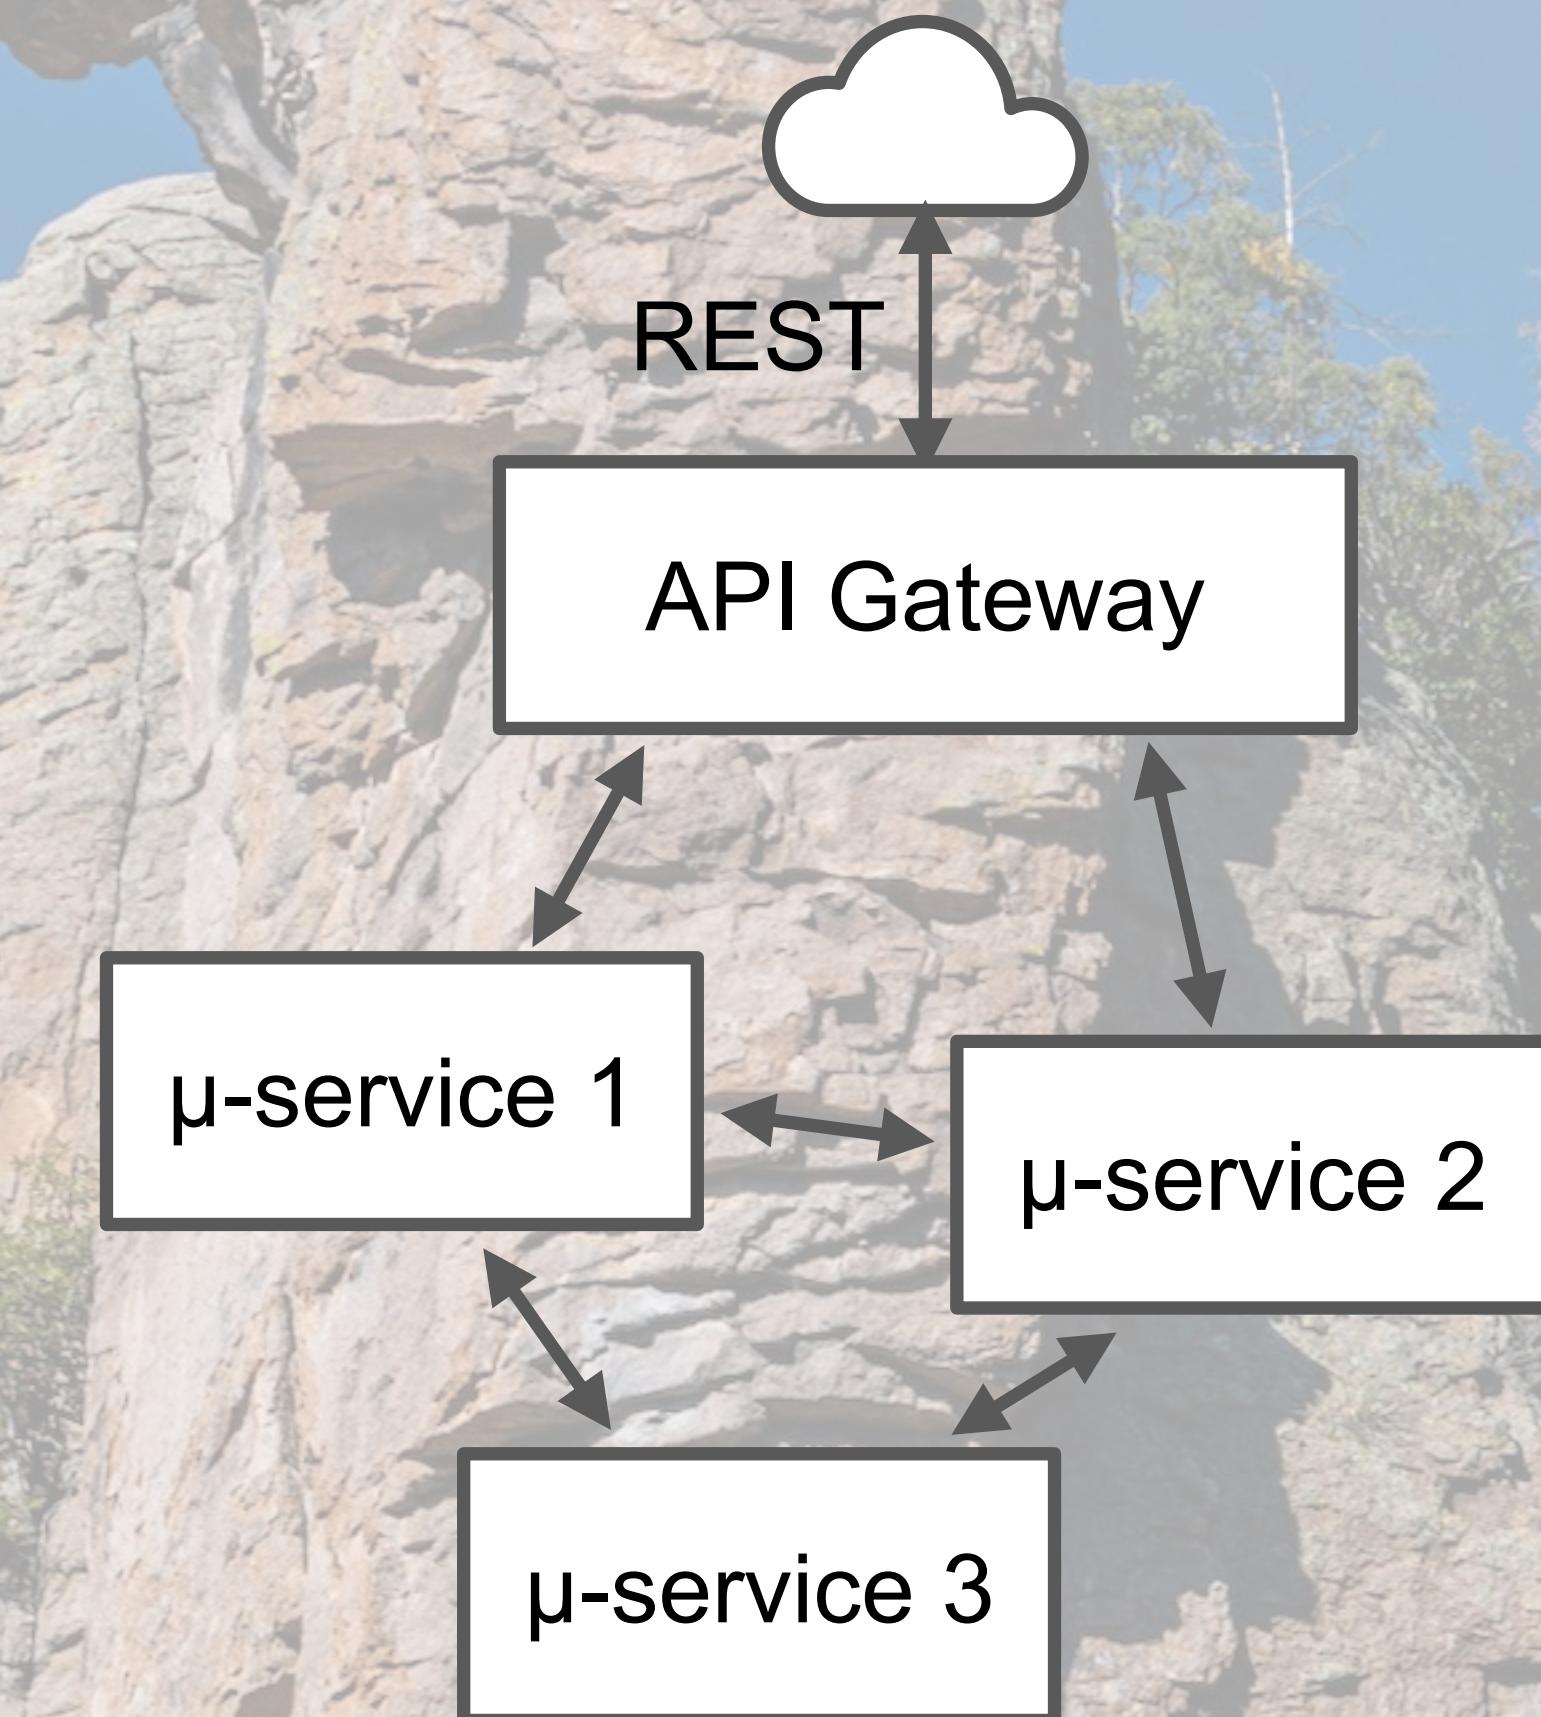
- They partition the domain
 - Conway's Law - Embraced
 - Separate responsibilities
- Separate management



What Are Microservices?

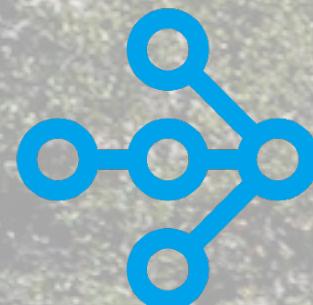
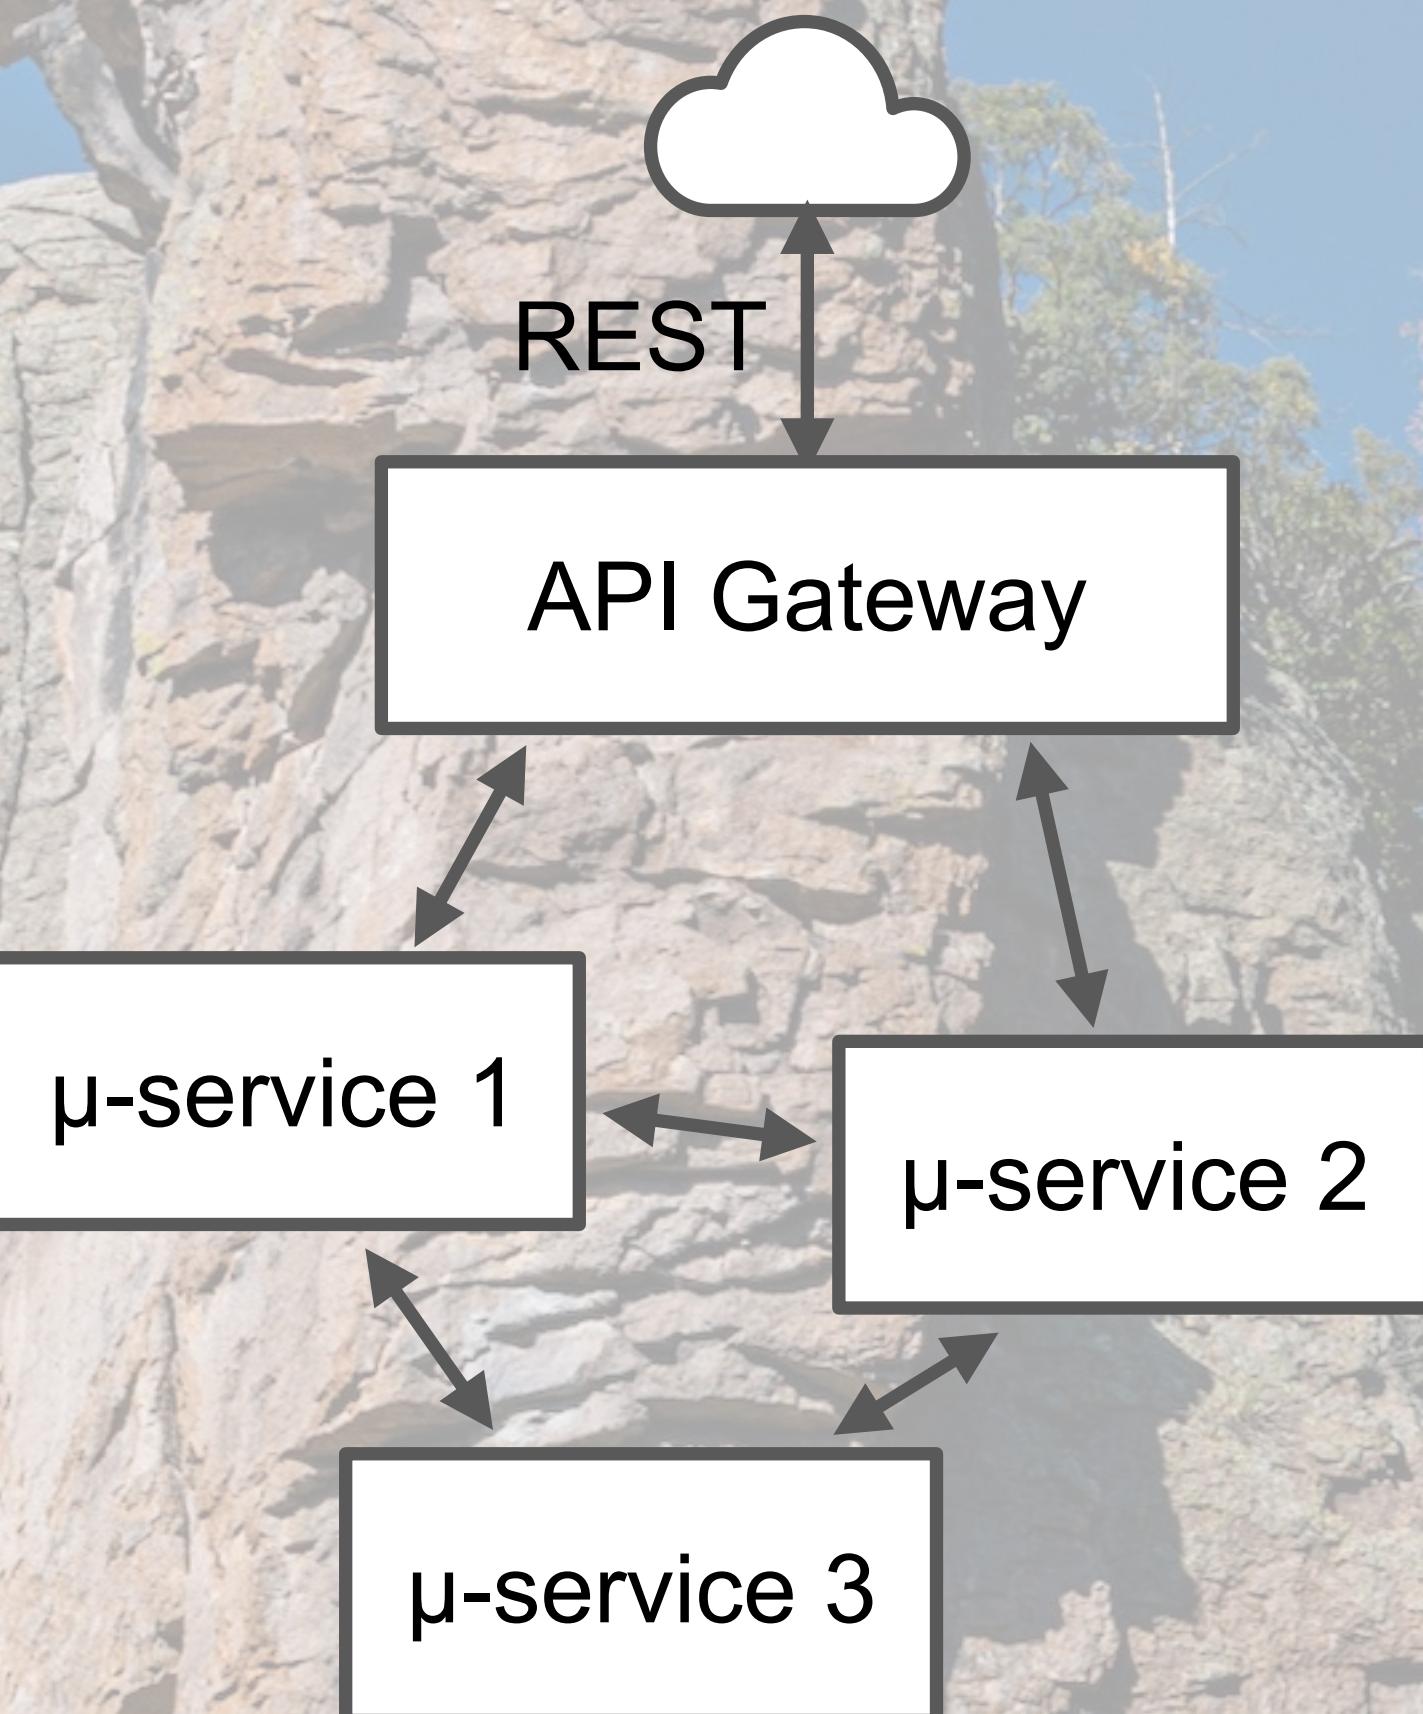
- They partition the domain
 - Conway's Law - Embraced
 - Separate responsibilities
- Separate management

What we mostly care
about for today's talk, the
“Ops in DevOps”



Conway's Law - Embraced

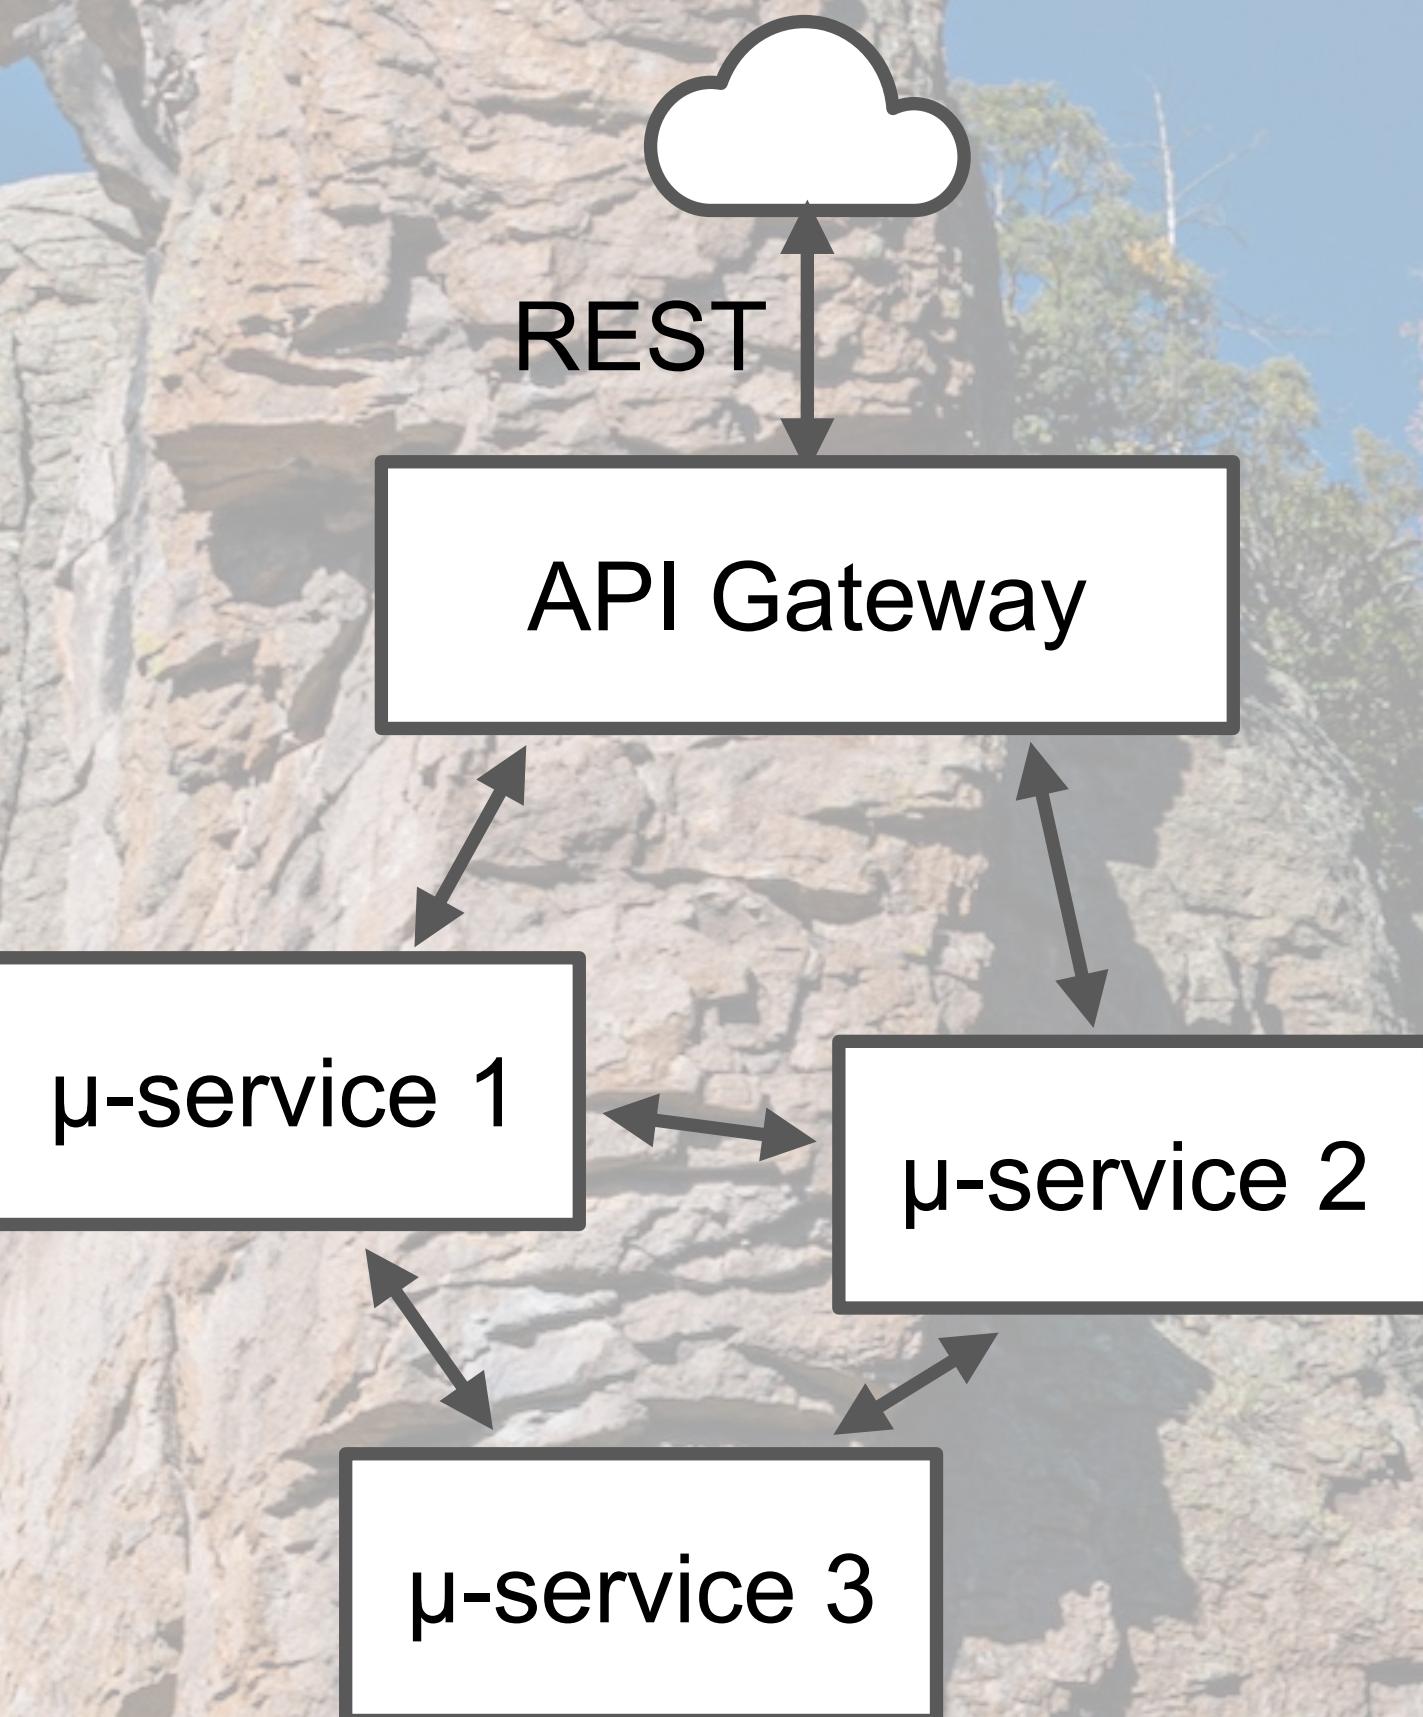
- “Any organization that designs a system will produce a design whose structure is a copy of the organization's communication structure”
- Let each team own and manage the services for its part of the domain



en.wikipedia.org/wiki/Conway's_law

Separate Responsibilities

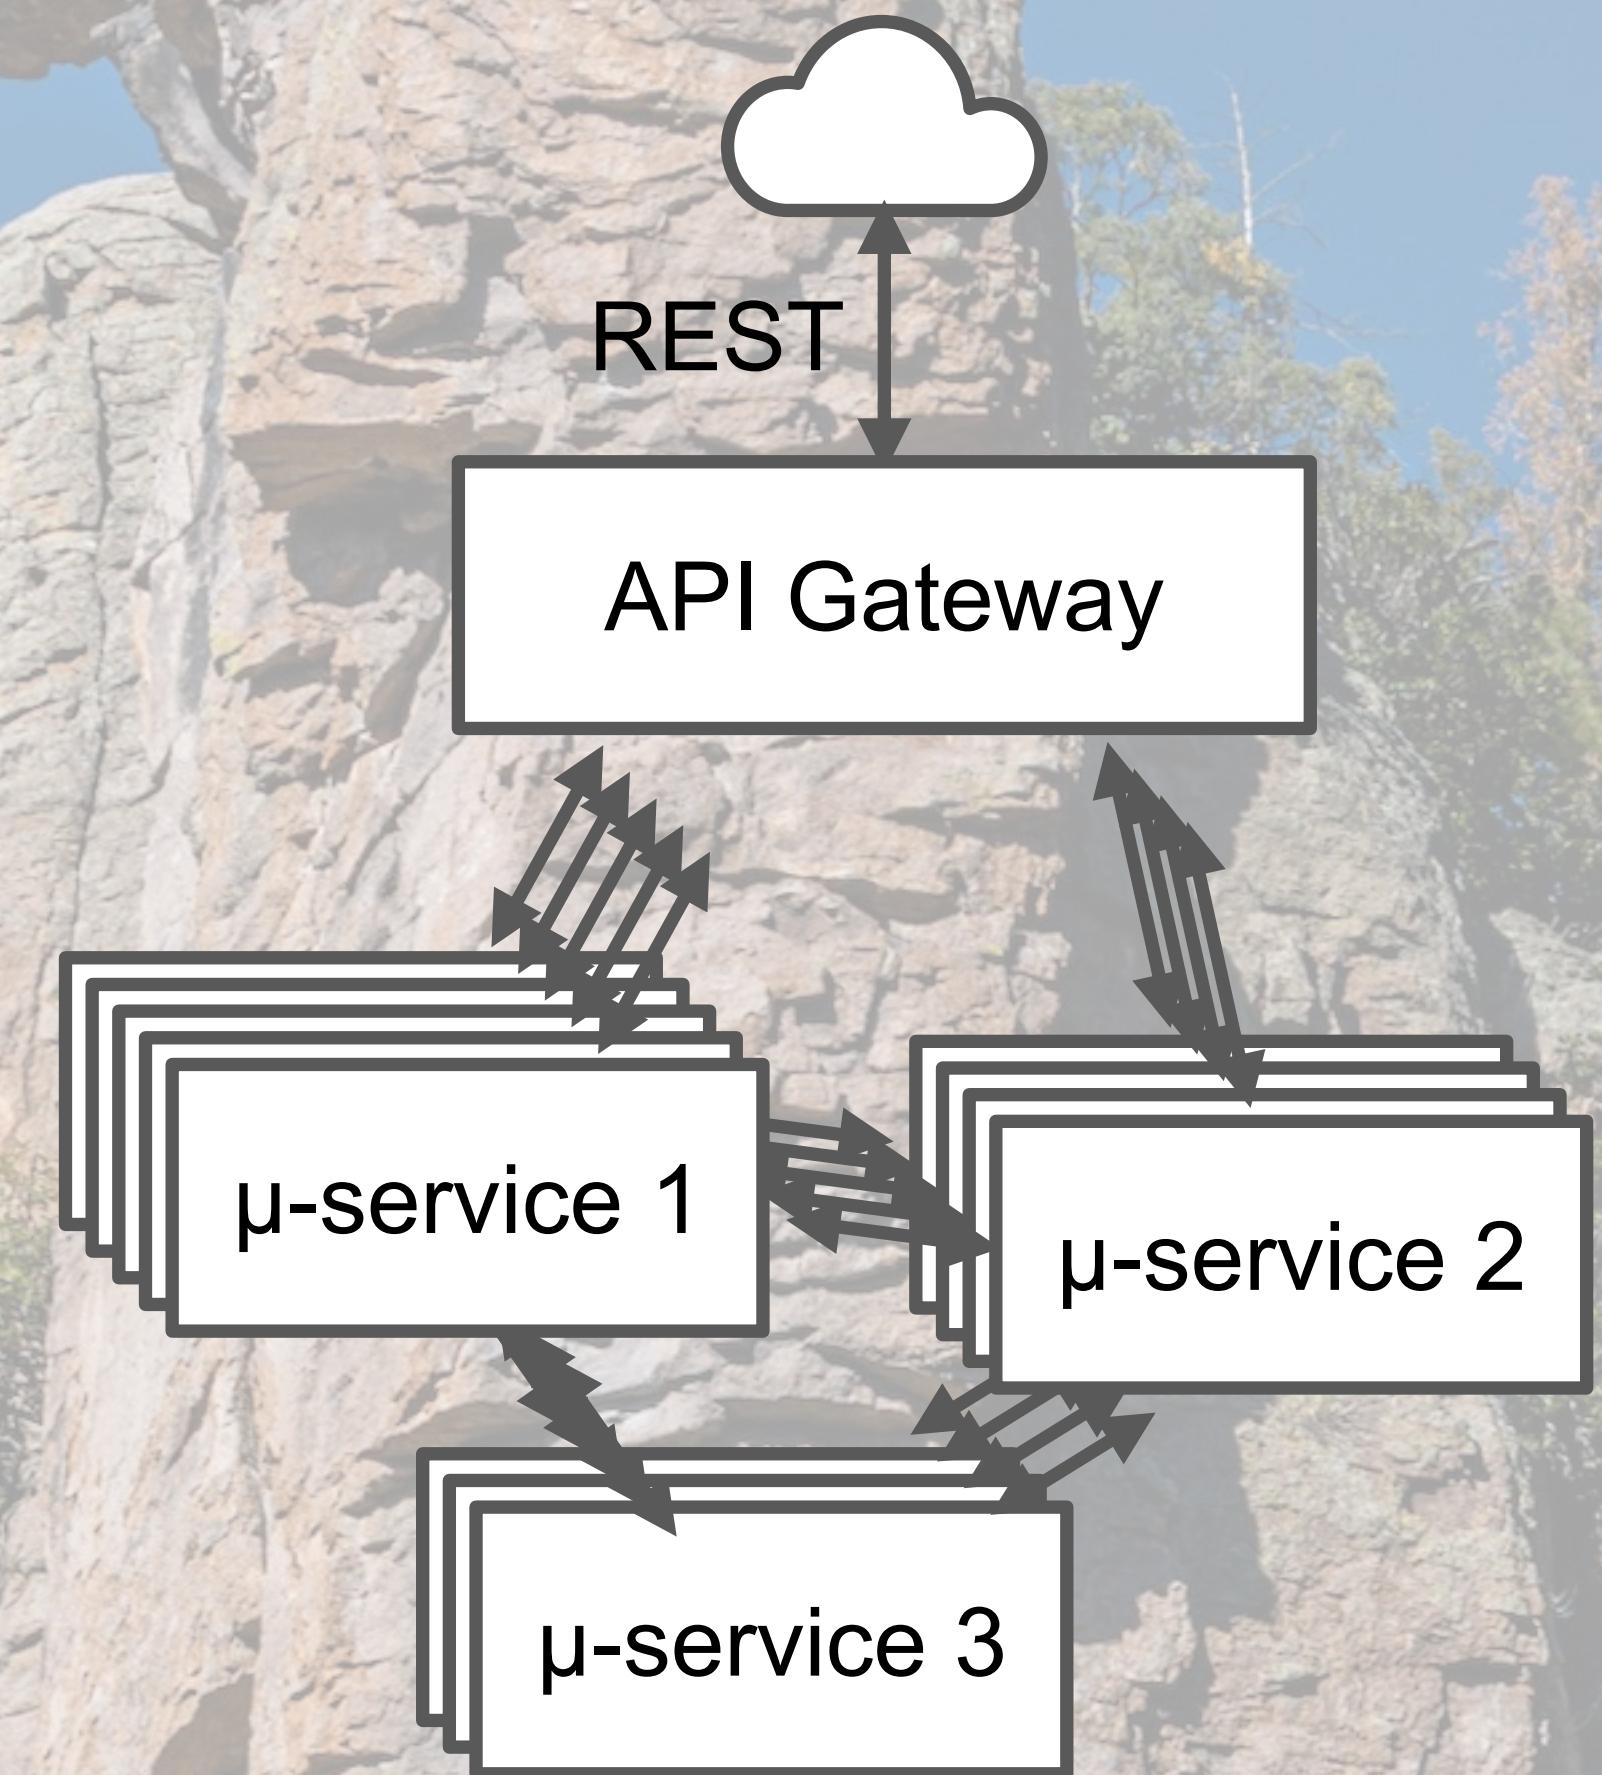
- Each microservice does “one thing”, a single responsibility with minimal coupling to the other microservices
- (Like, hopefully, the teams are organized, too...)



wikipedia.org/wiki/Single-responsibility_principle

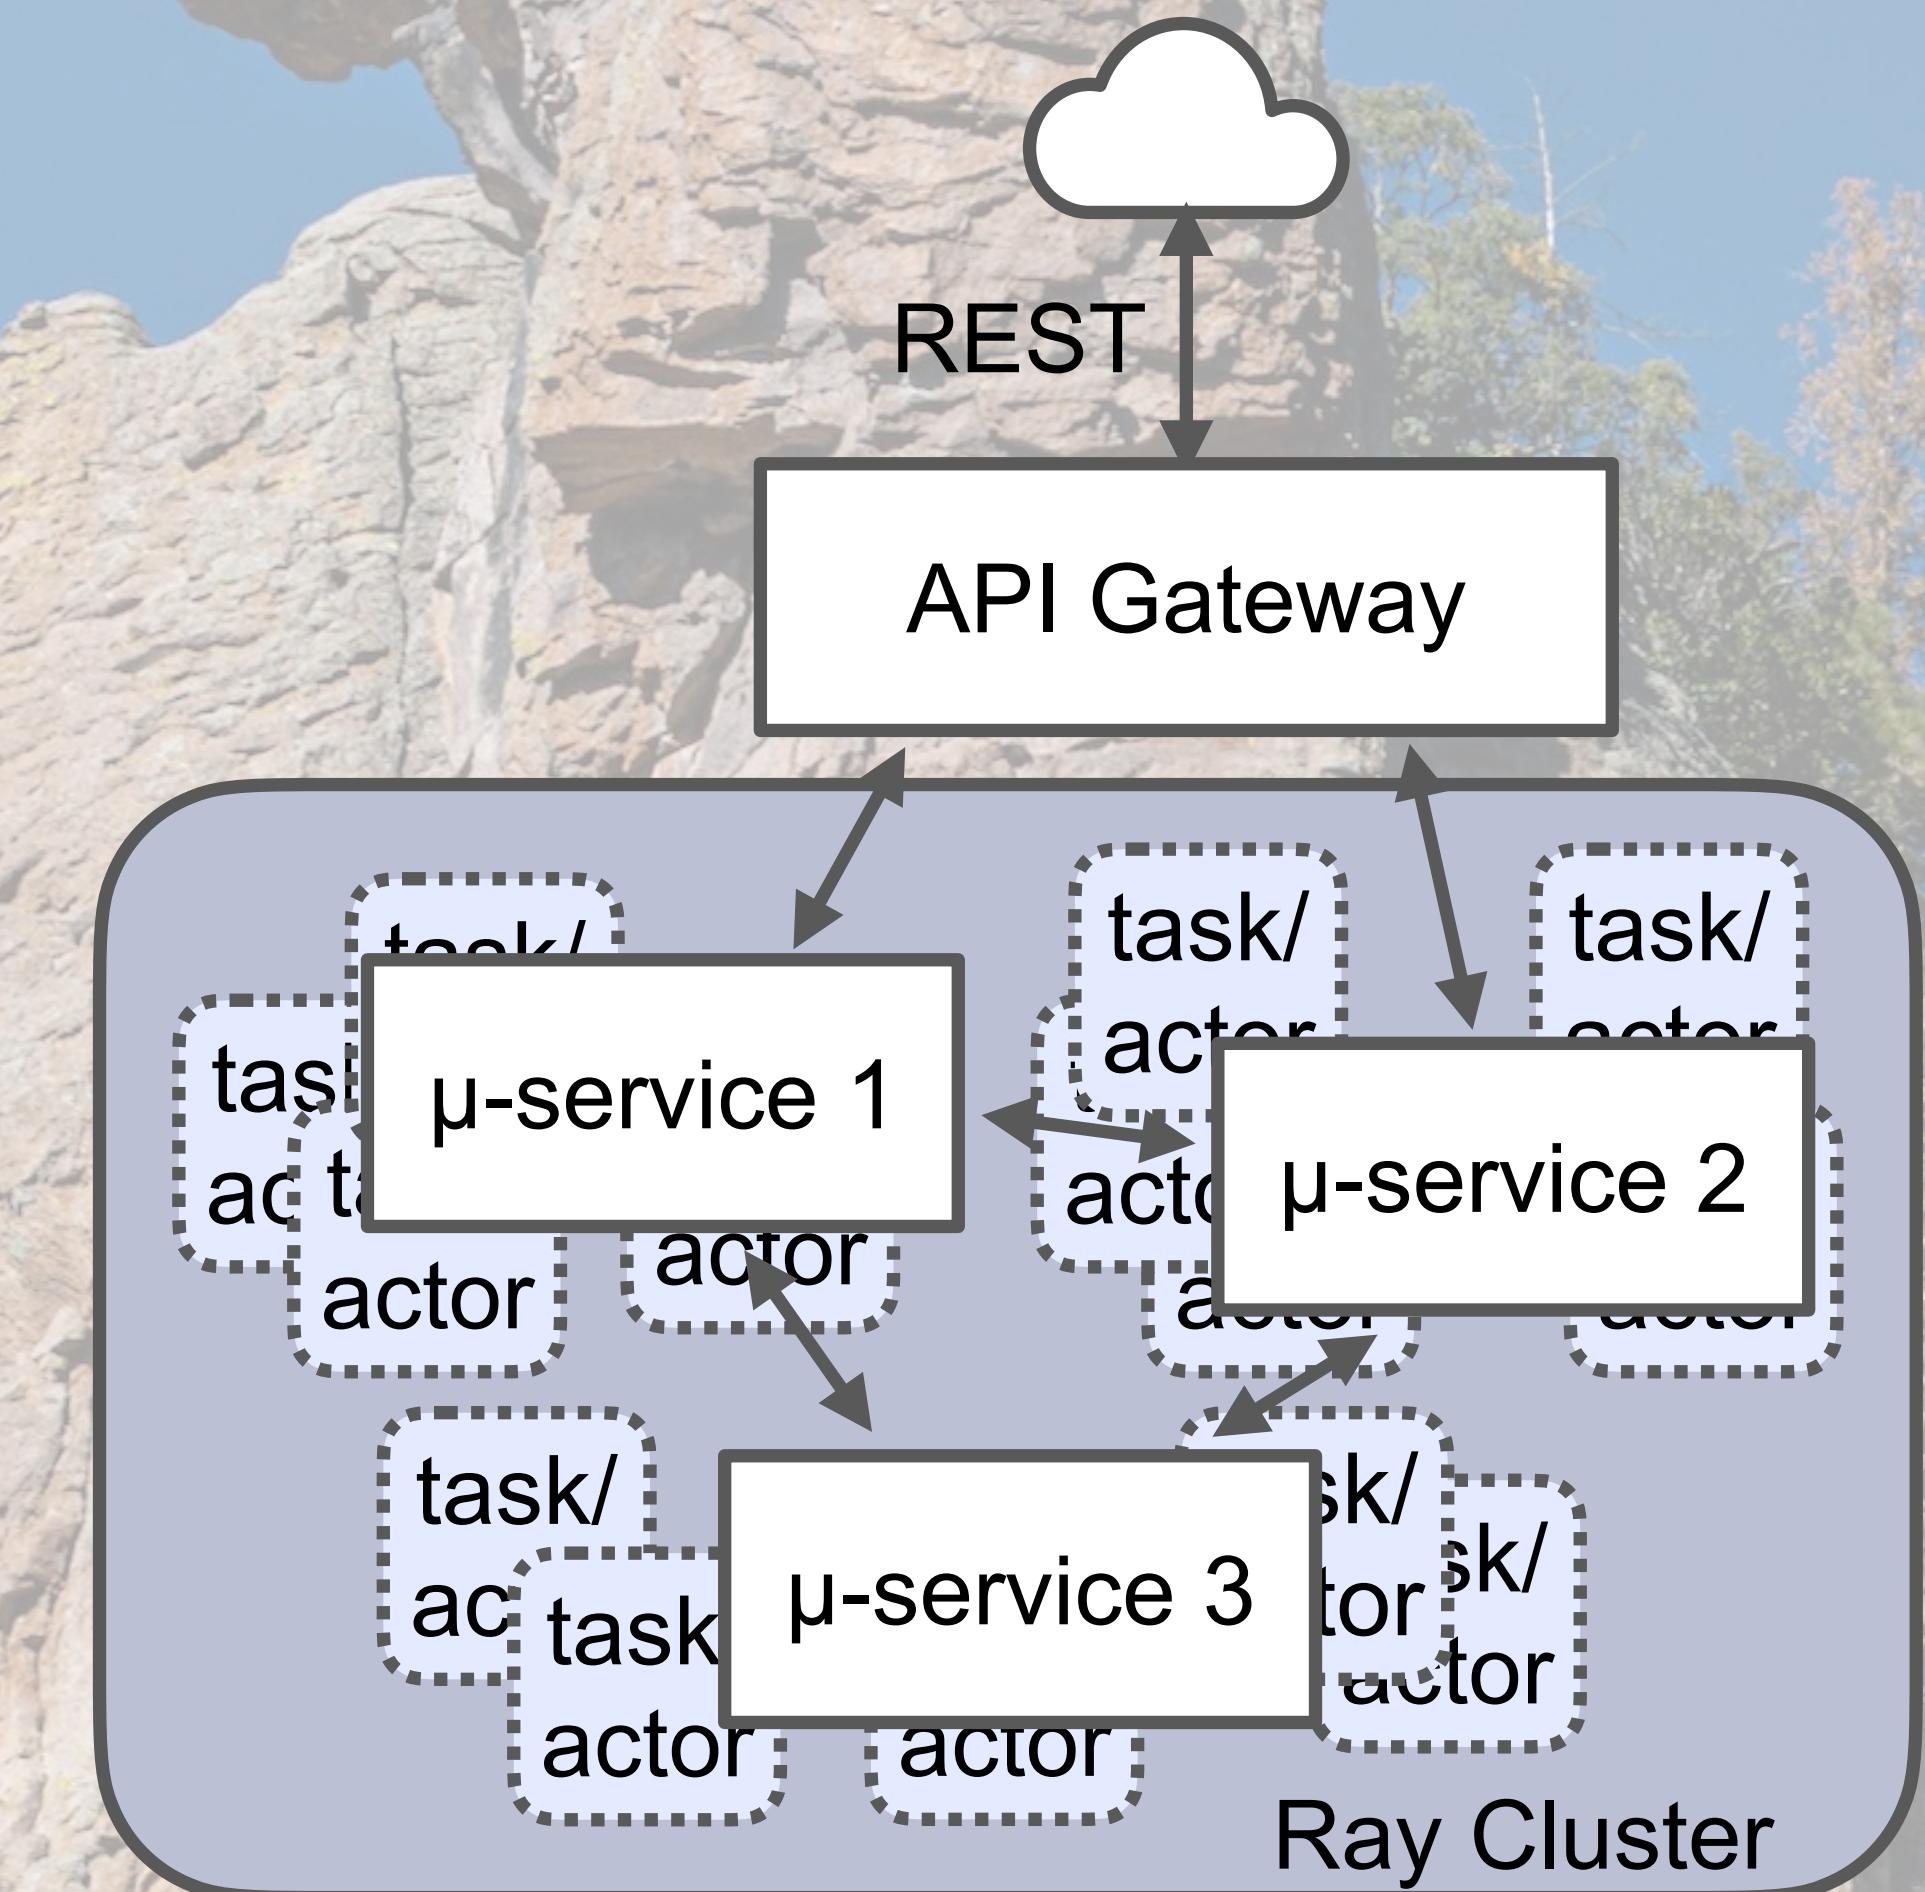
Separate Management

- Each team manages its own instances
- Each microservice has a different number of instances for scalability and resiliency
- But they have to be managed **explicitly**



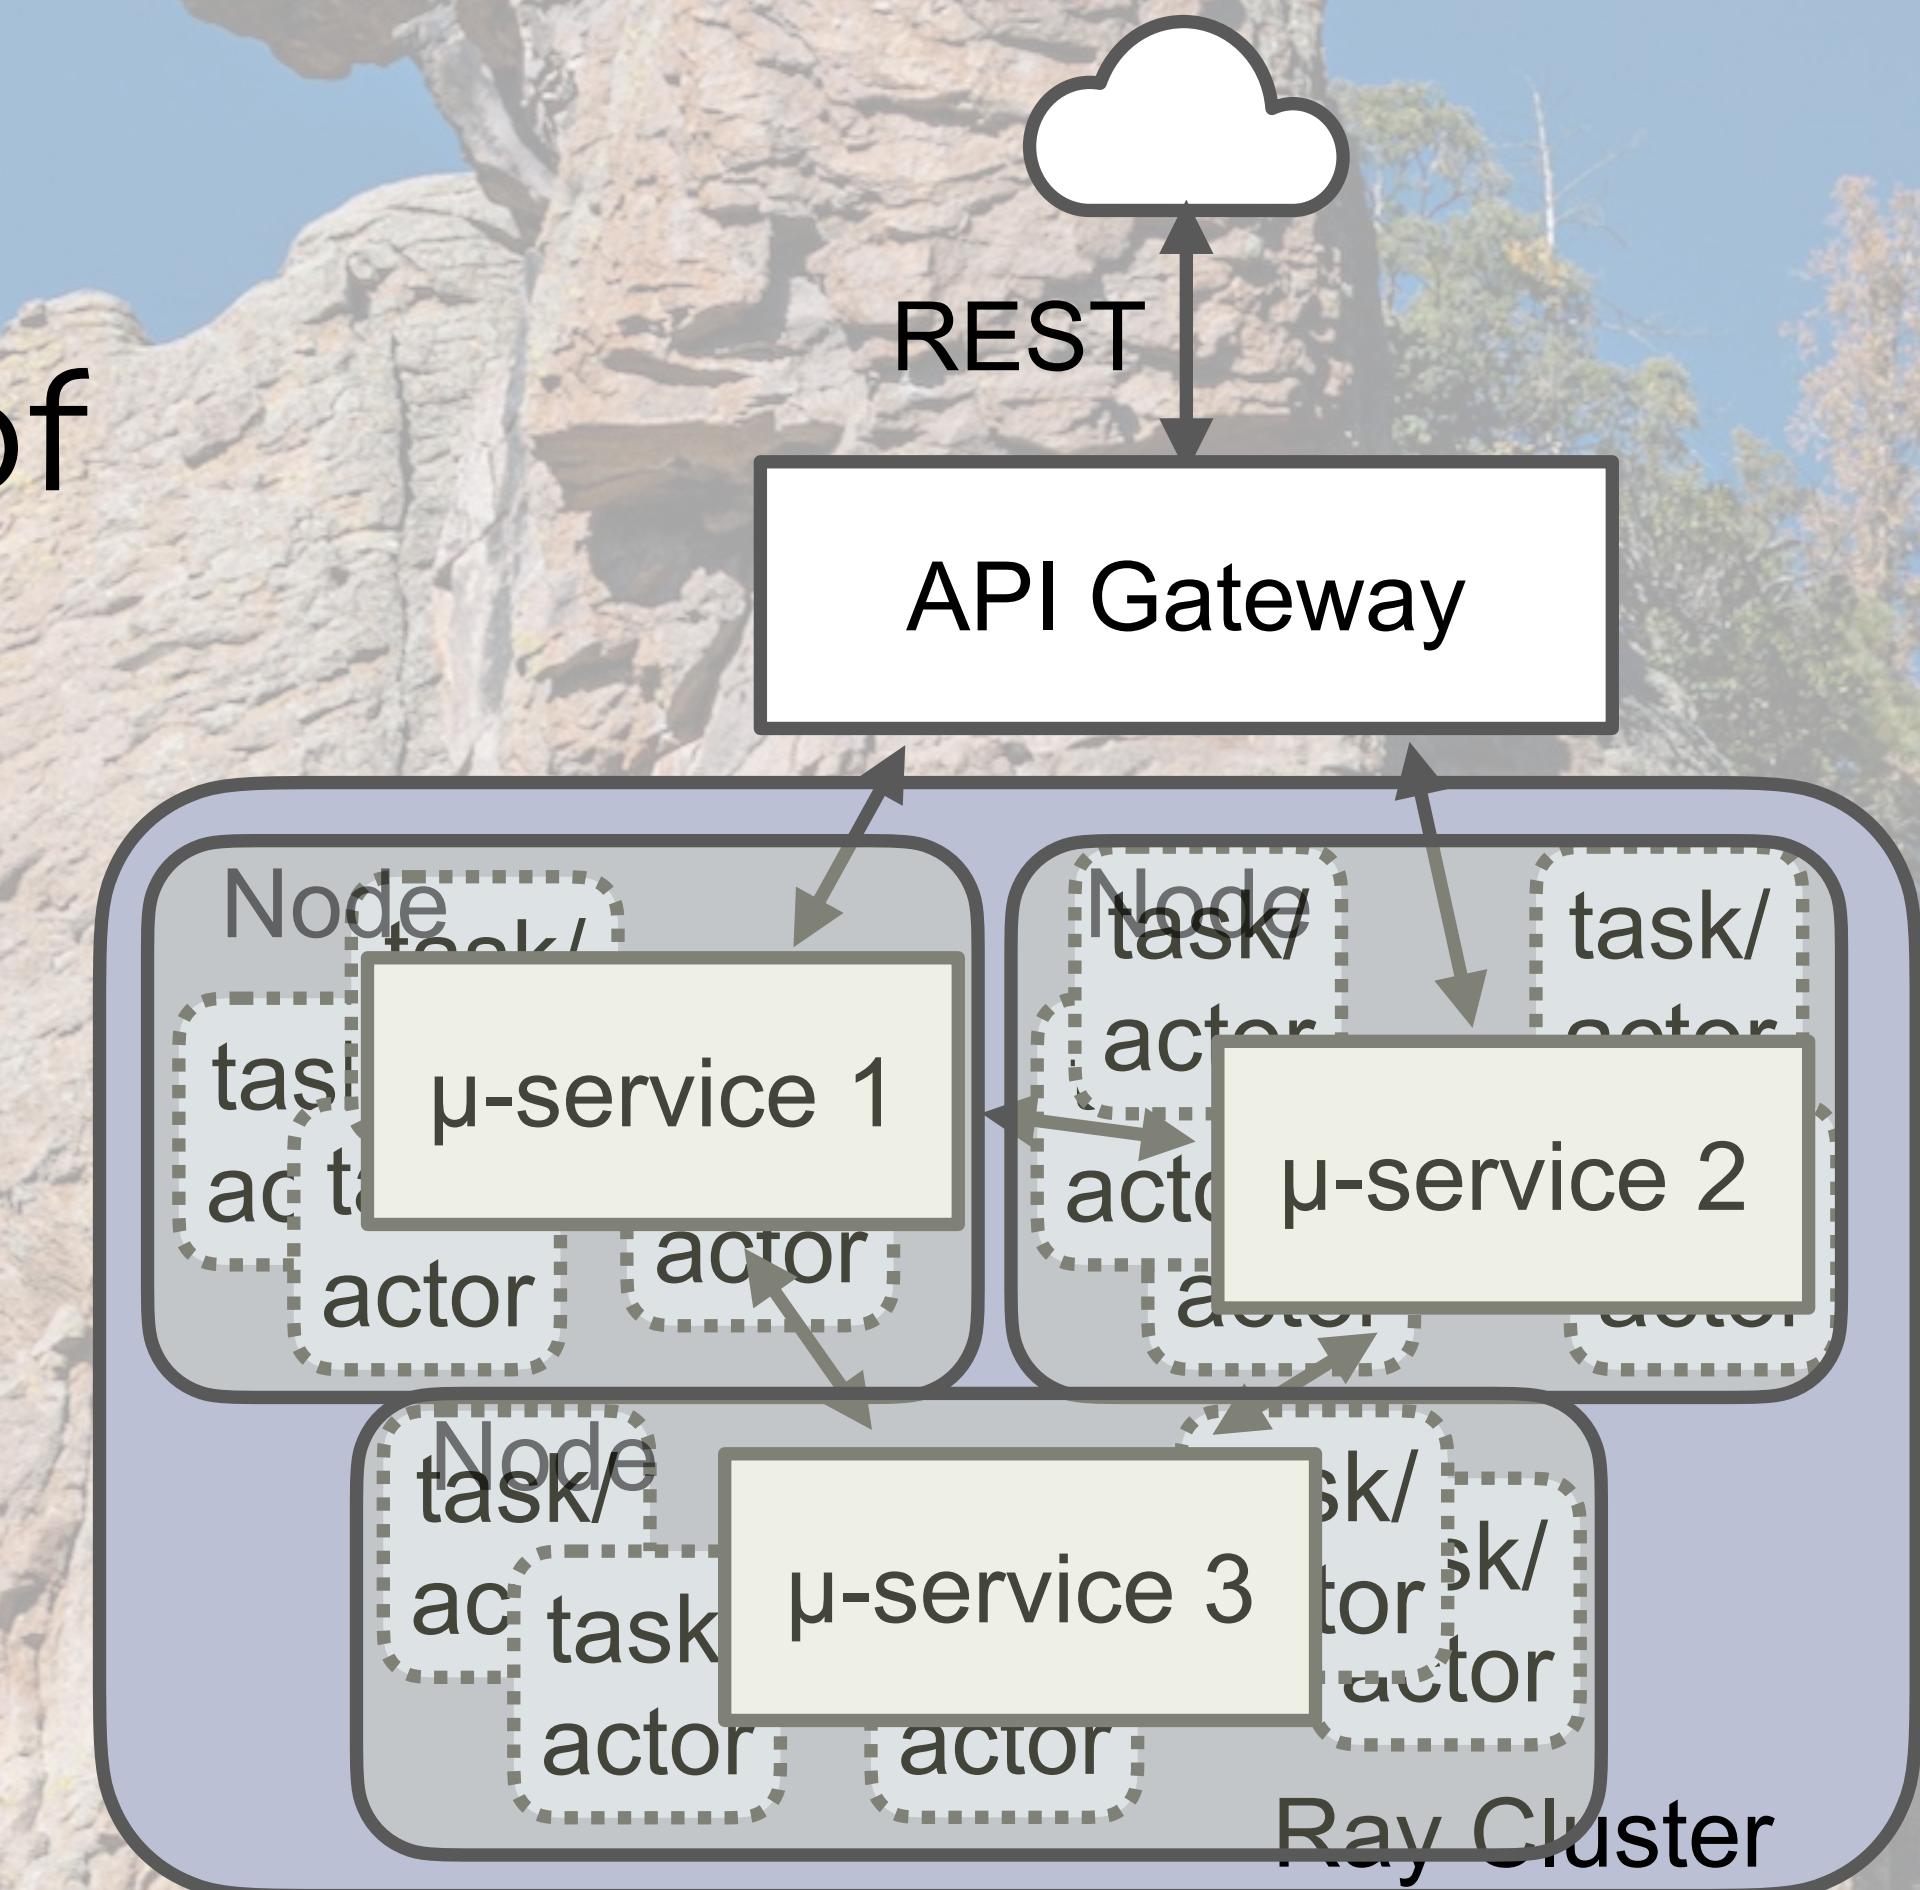
Management - Simplified

- With Ray, you have one “logical” instance to manage and Ray does the cluster-wide scaling for you.



What about Kubernetes (and others...)?

- Ray scaling is very fine grained.
- It operates within the “nodes” of coarse-grained managers
- Containers, pods, VMs, or physical machines





Adopting Ray and the Ray community





@deanwampler

If you're already using...

- joblib
- multiprocessing.Pool

For example, from this:

```
from multiprocessing.pool import Pool
```

To this:

```
from ray.util.multiprocessing.pool import Pool
```

- Use Ray's implementations
 - Drop-in replacements
 - Change import statements
 - Break the one-node limitation!

- ... And Ray is integrated with asyncio

See these blog posts:

<https://medium.com/distributed-computing-with-ray/how-to-scale-python-multiprocessing-to-a-cluster-with-one-line-of-code-d19f242f60ff>

<https://medium.com/distributed-computing-with-ray/easy-distributed-scikit-learn-training-with-ray-54ff8b643b33>



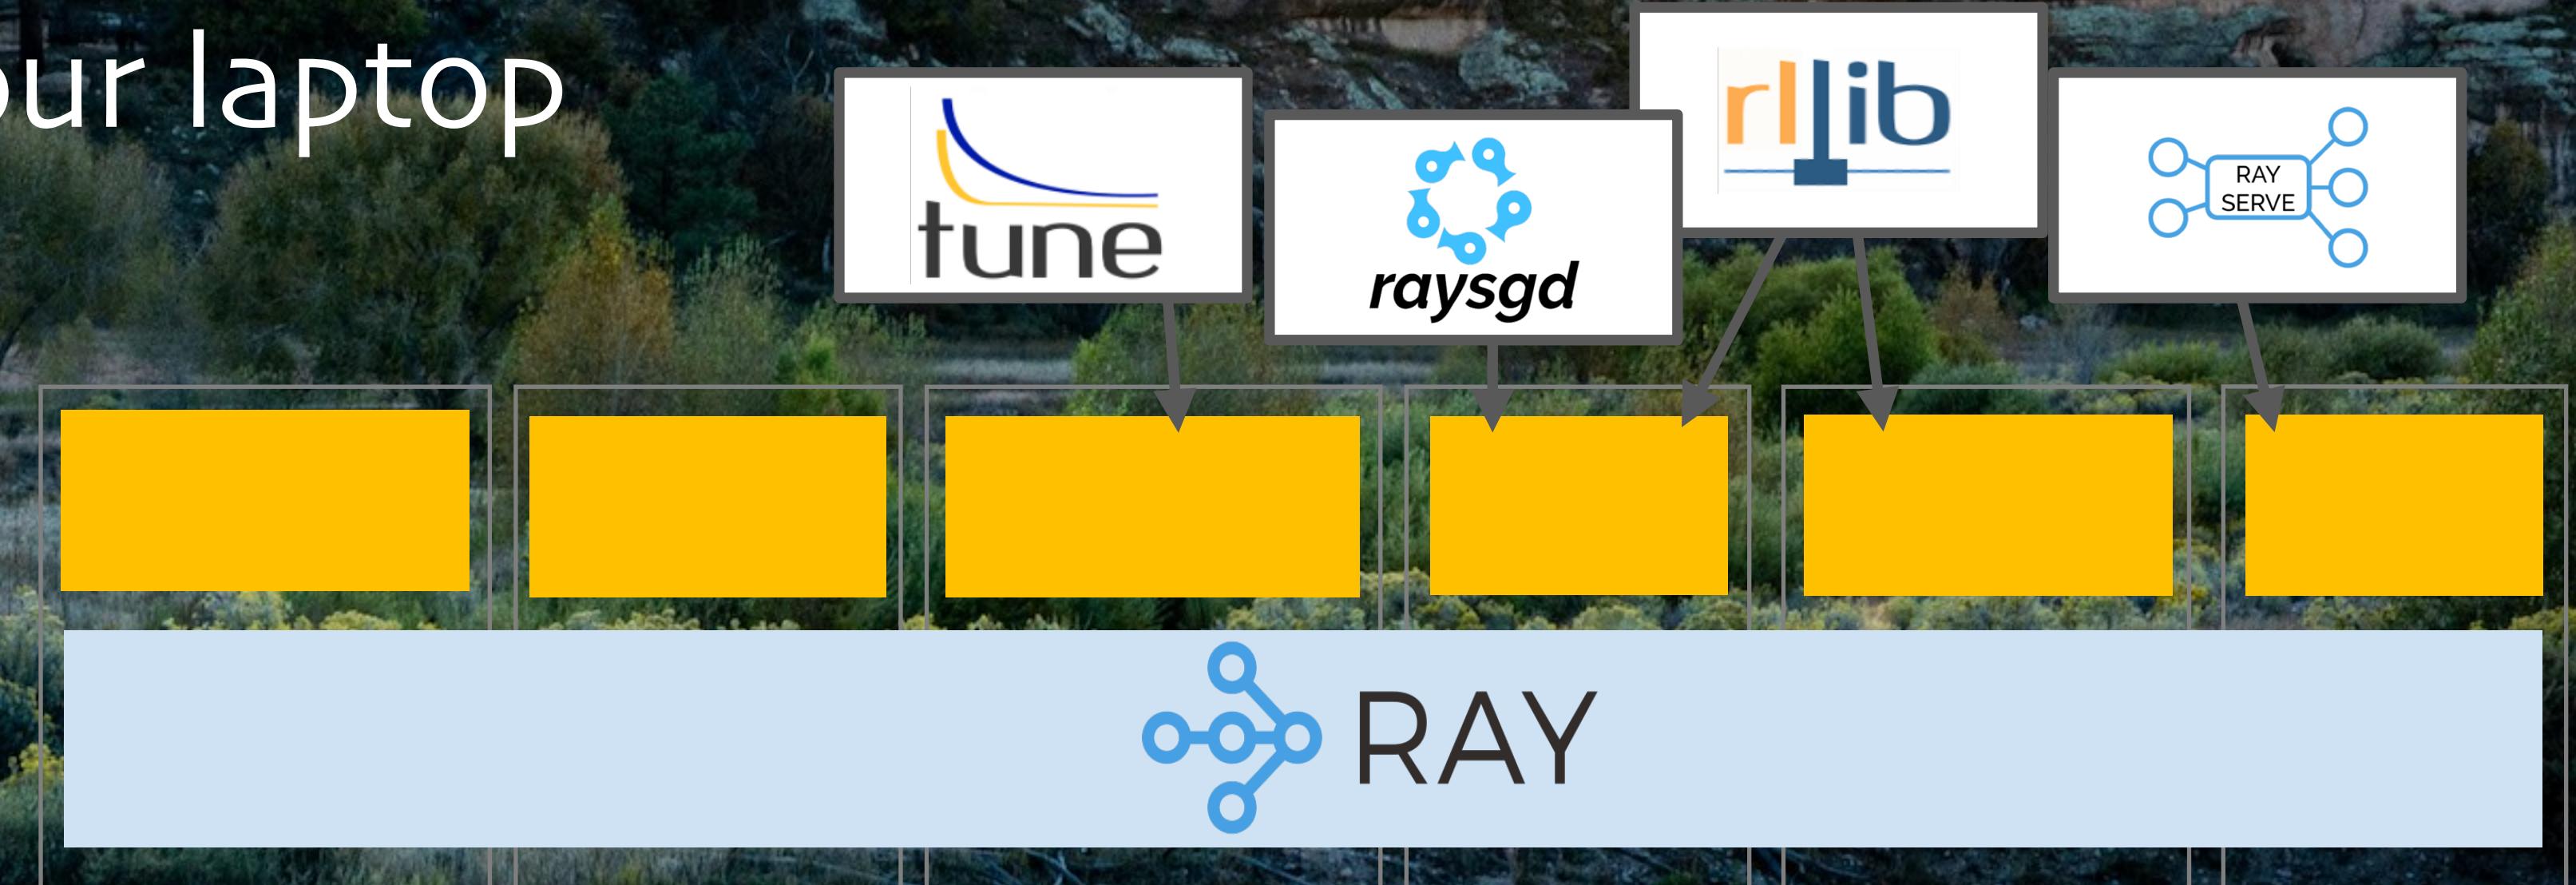
Ray Community and Resources

- ray.io
- Tutorials (free): anyscale.com/academy
- Need help?
 - Ray Slack: ray-distributed.slack.com
 - [ray-dev](https://groups.google.com/g/ray-dev) Google group



Conclusion

- Ray is the new state-of-the-art for distributed computing
 - The shortest path from your laptop to the cloud
 - Run complex distributed tasks on large clusters from simple code on your laptop





ray.io

dean@deanwampler.com

@deanwampler

dominodatalab.com

Slides at
polyglotprogramming.com/talks