

Scaling ML/AI Applications with Ray

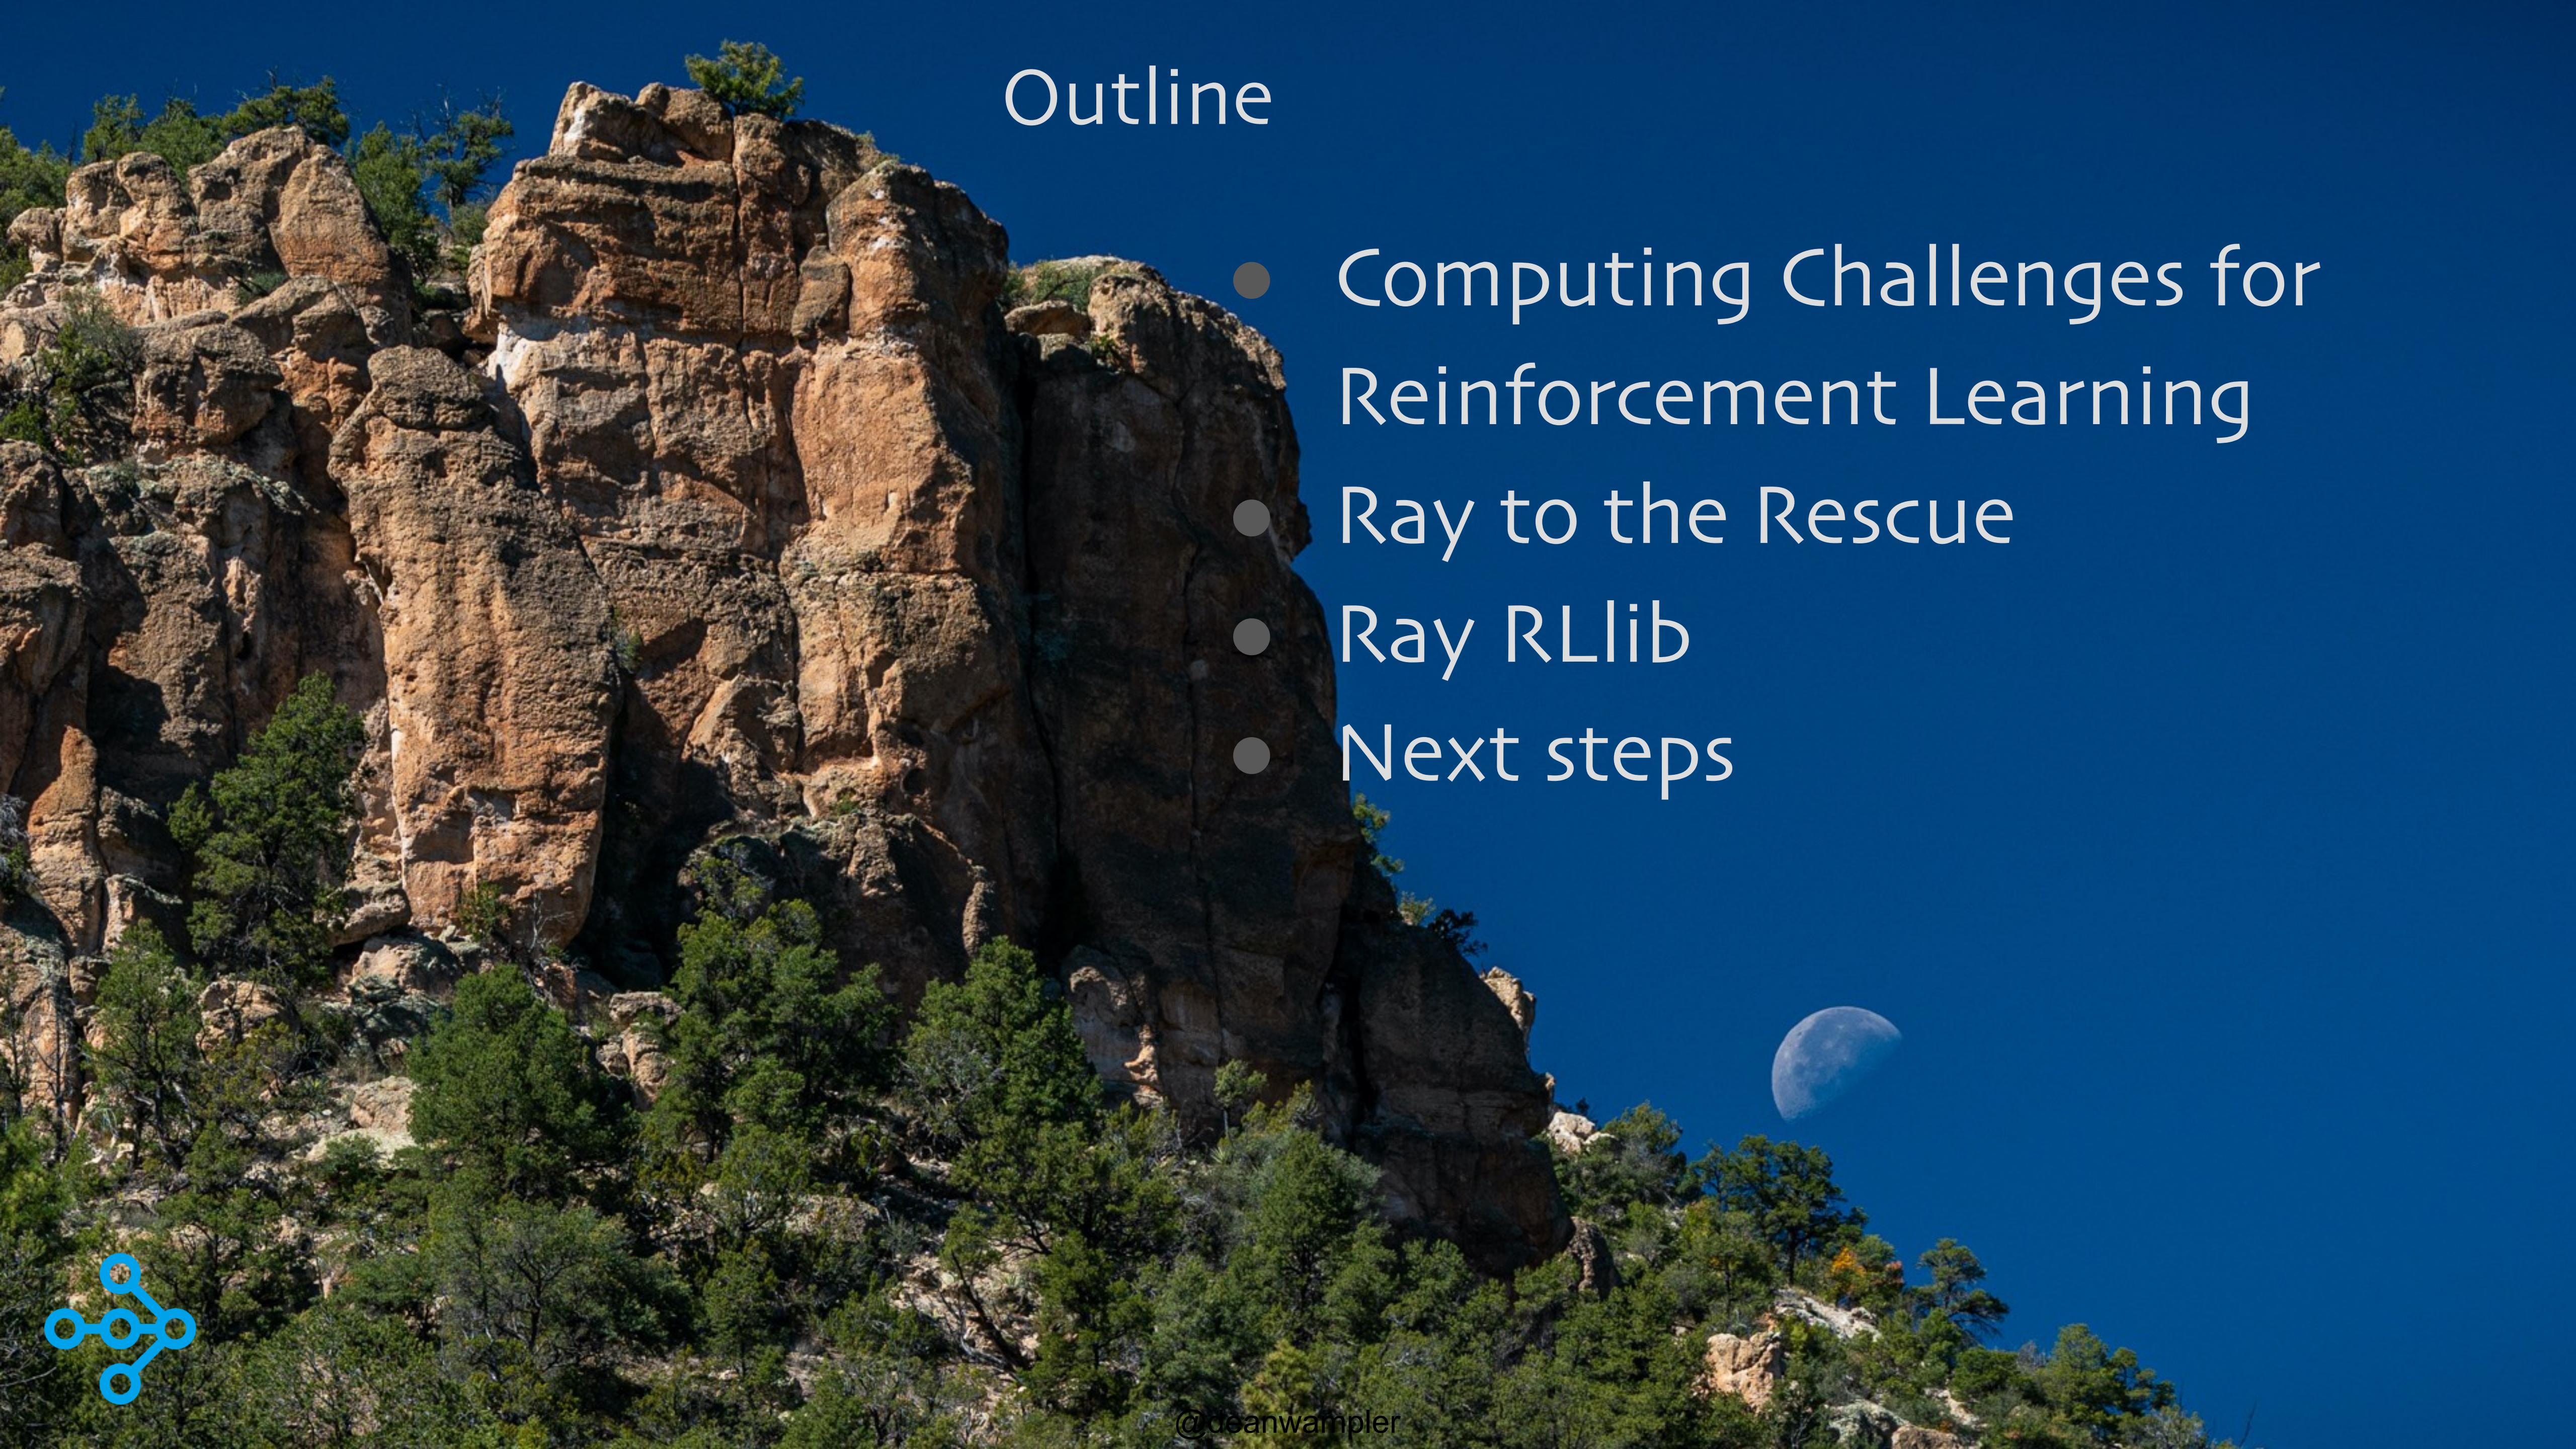
Dean Wampler

October 5, 2020

dean@deanwampler.com

[@deanwampler](https://twitter.com/deanwampler)

ray.io



Outline

- Computing Challenges for Reinforcement Learning
- Ray to the Rescue
- Ray RLlib
- Next steps



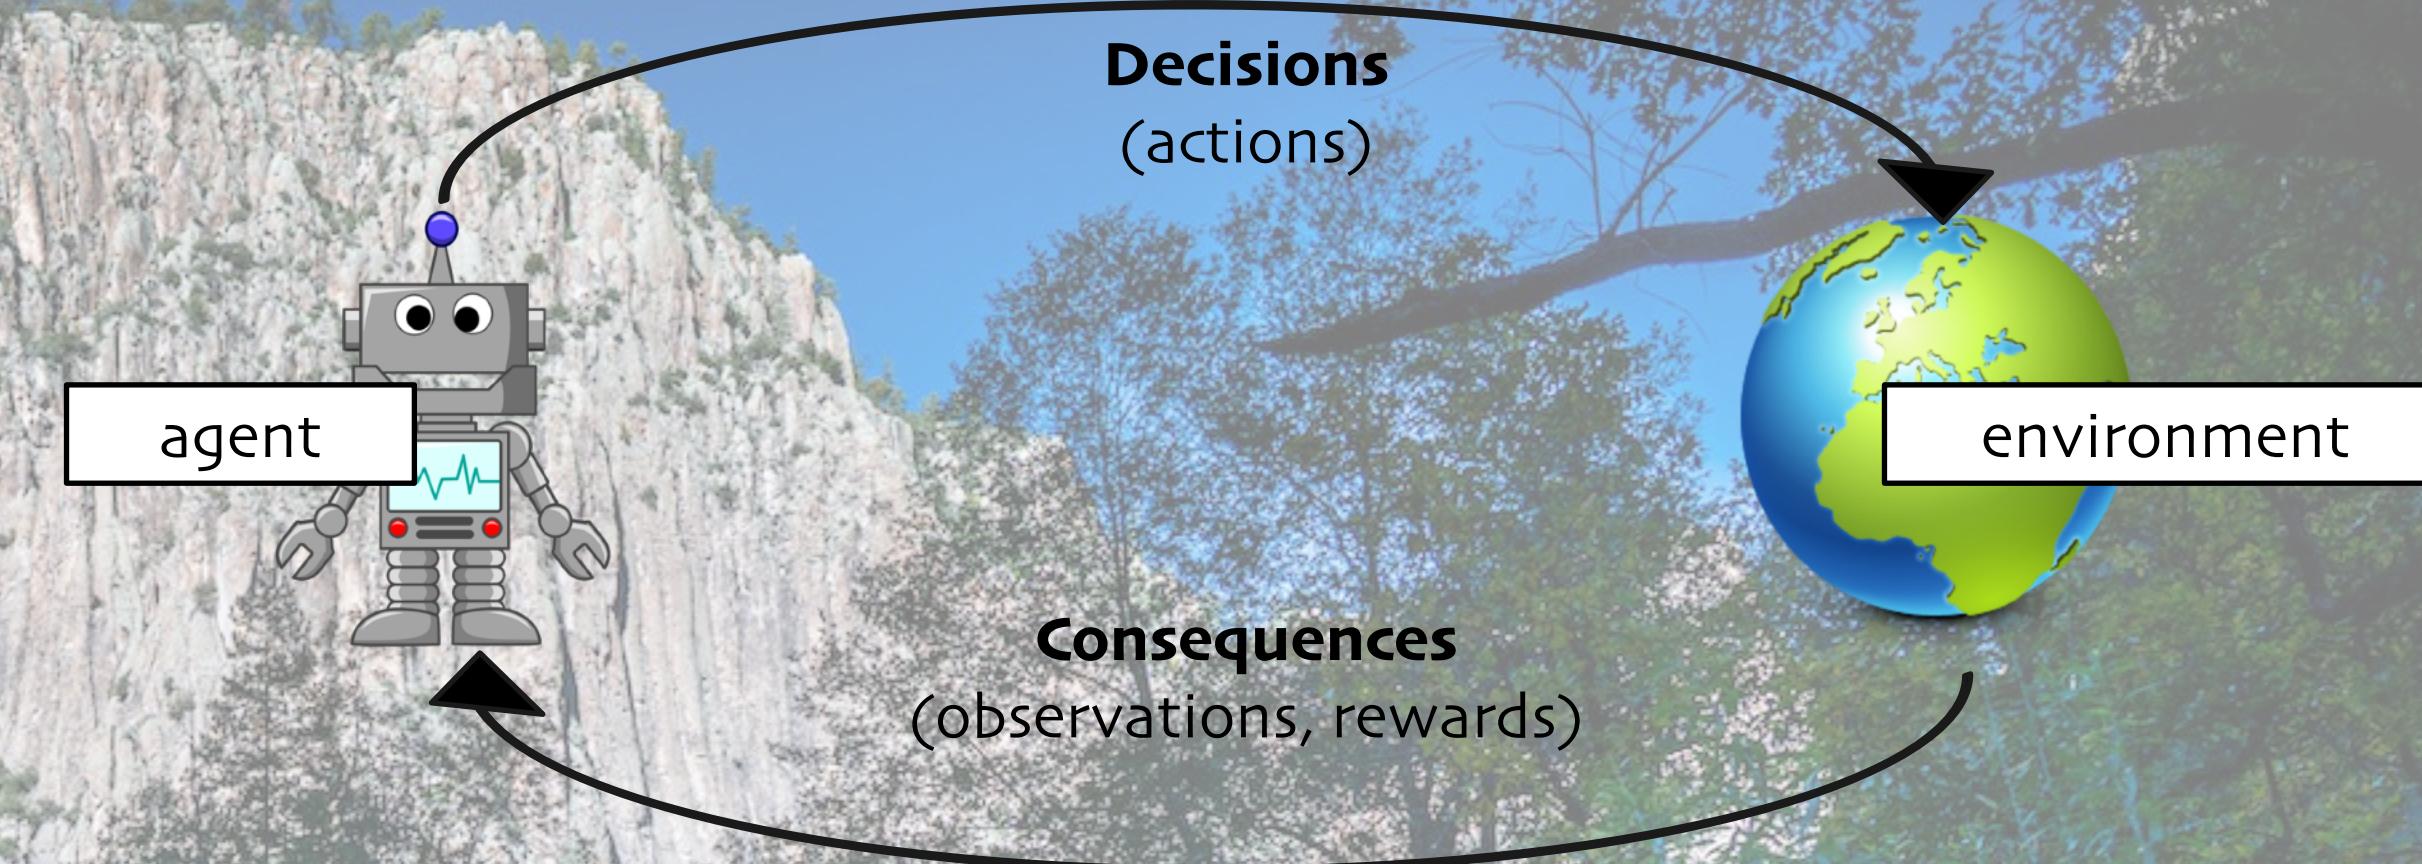


Computing Challenges for Reinforcement Learning



@deanwampler

Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

Industrial
Processes

System
Optimization

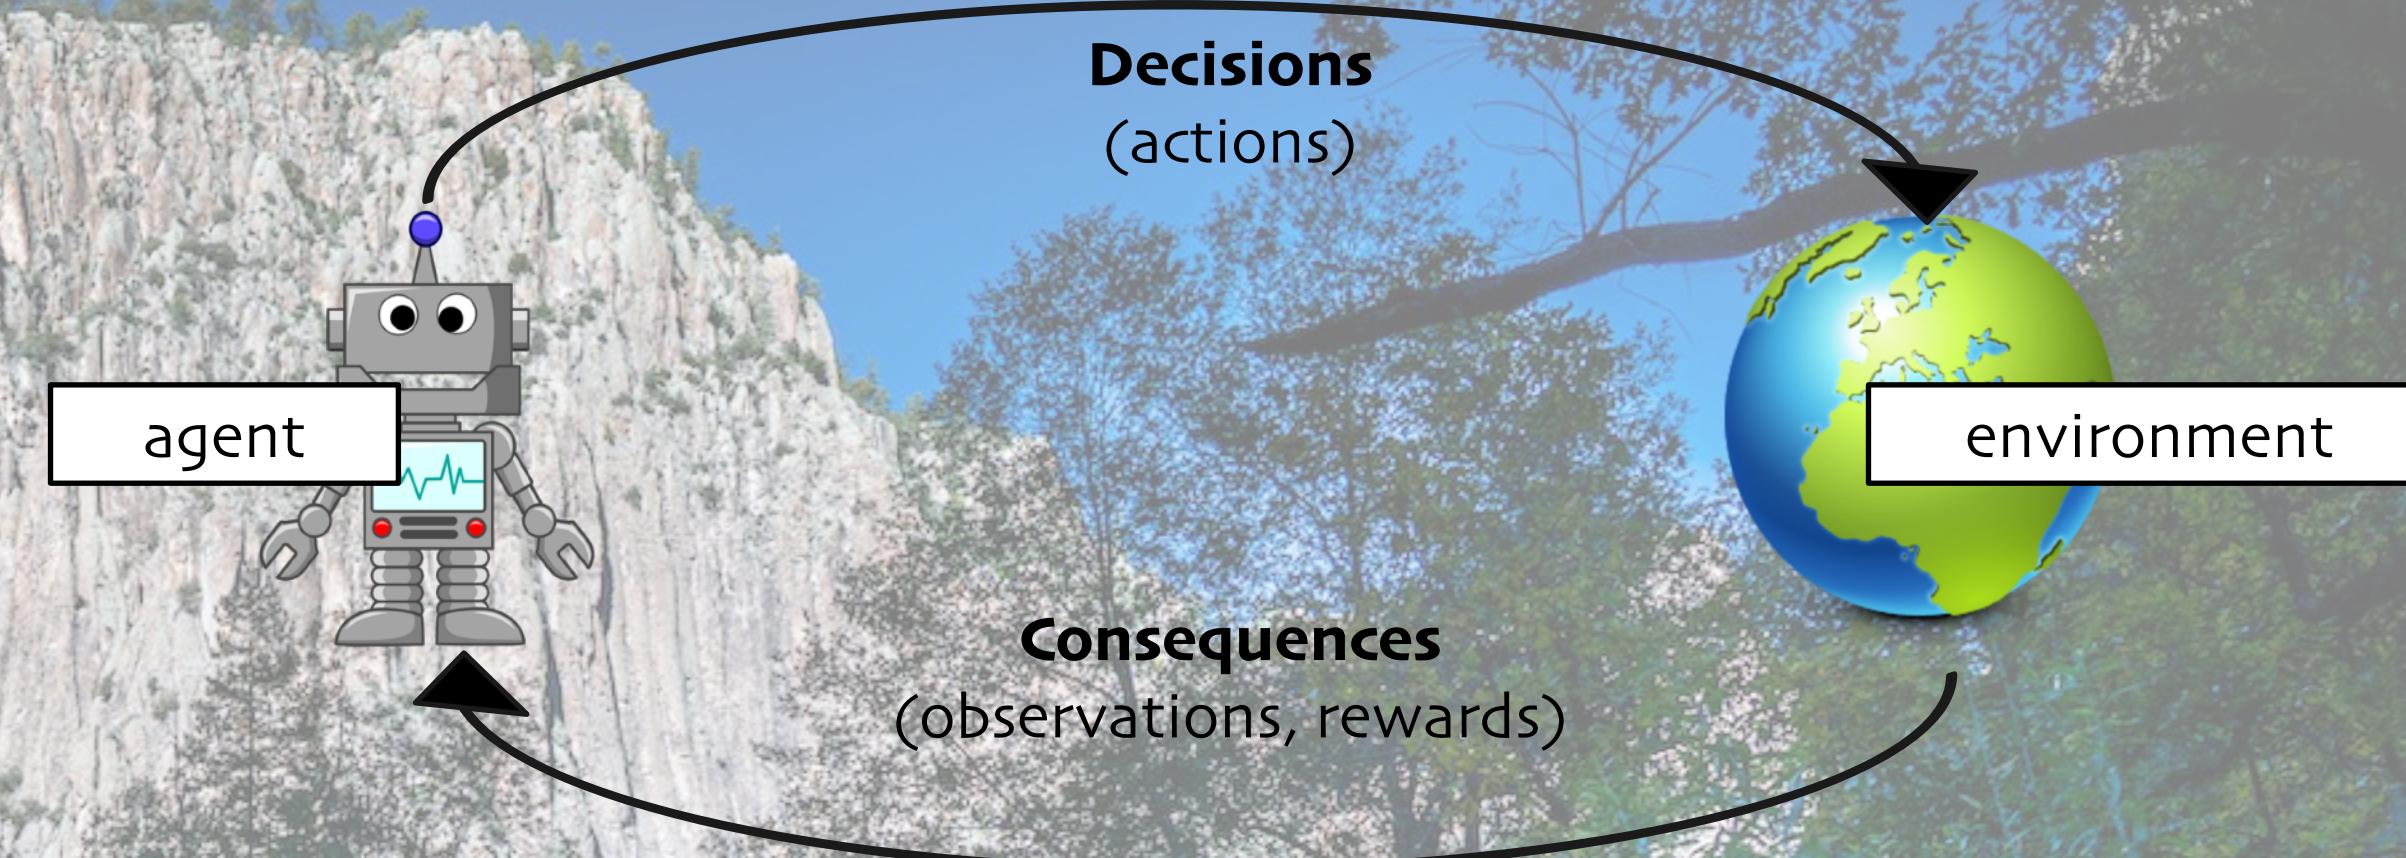
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

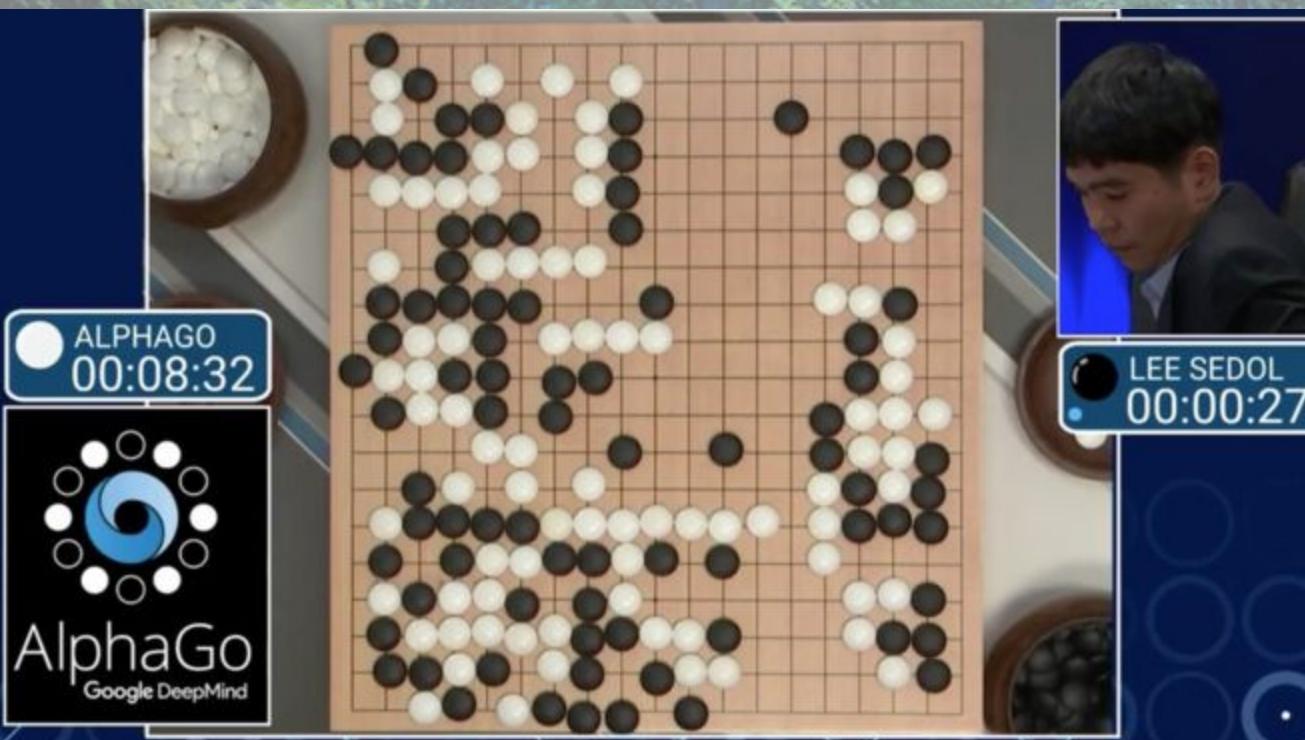
Industrial
Processes

System
Optimization

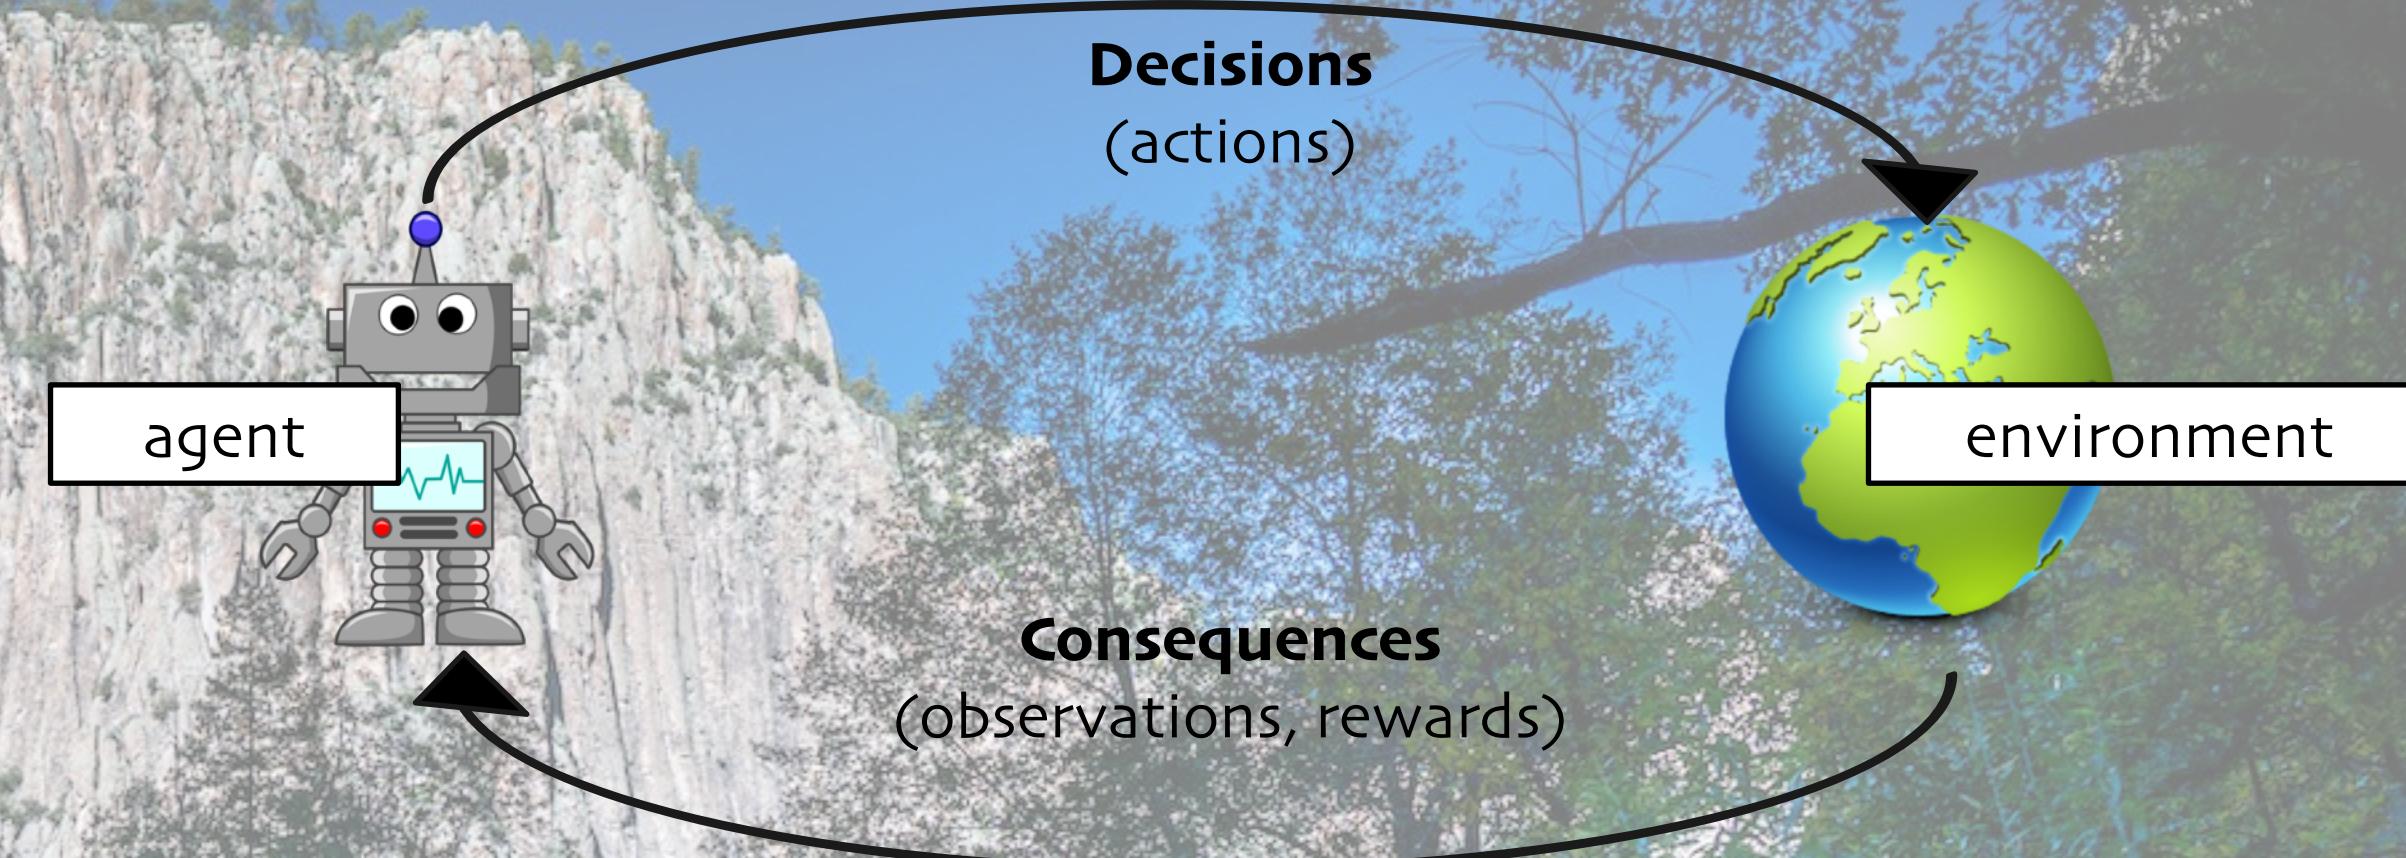
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

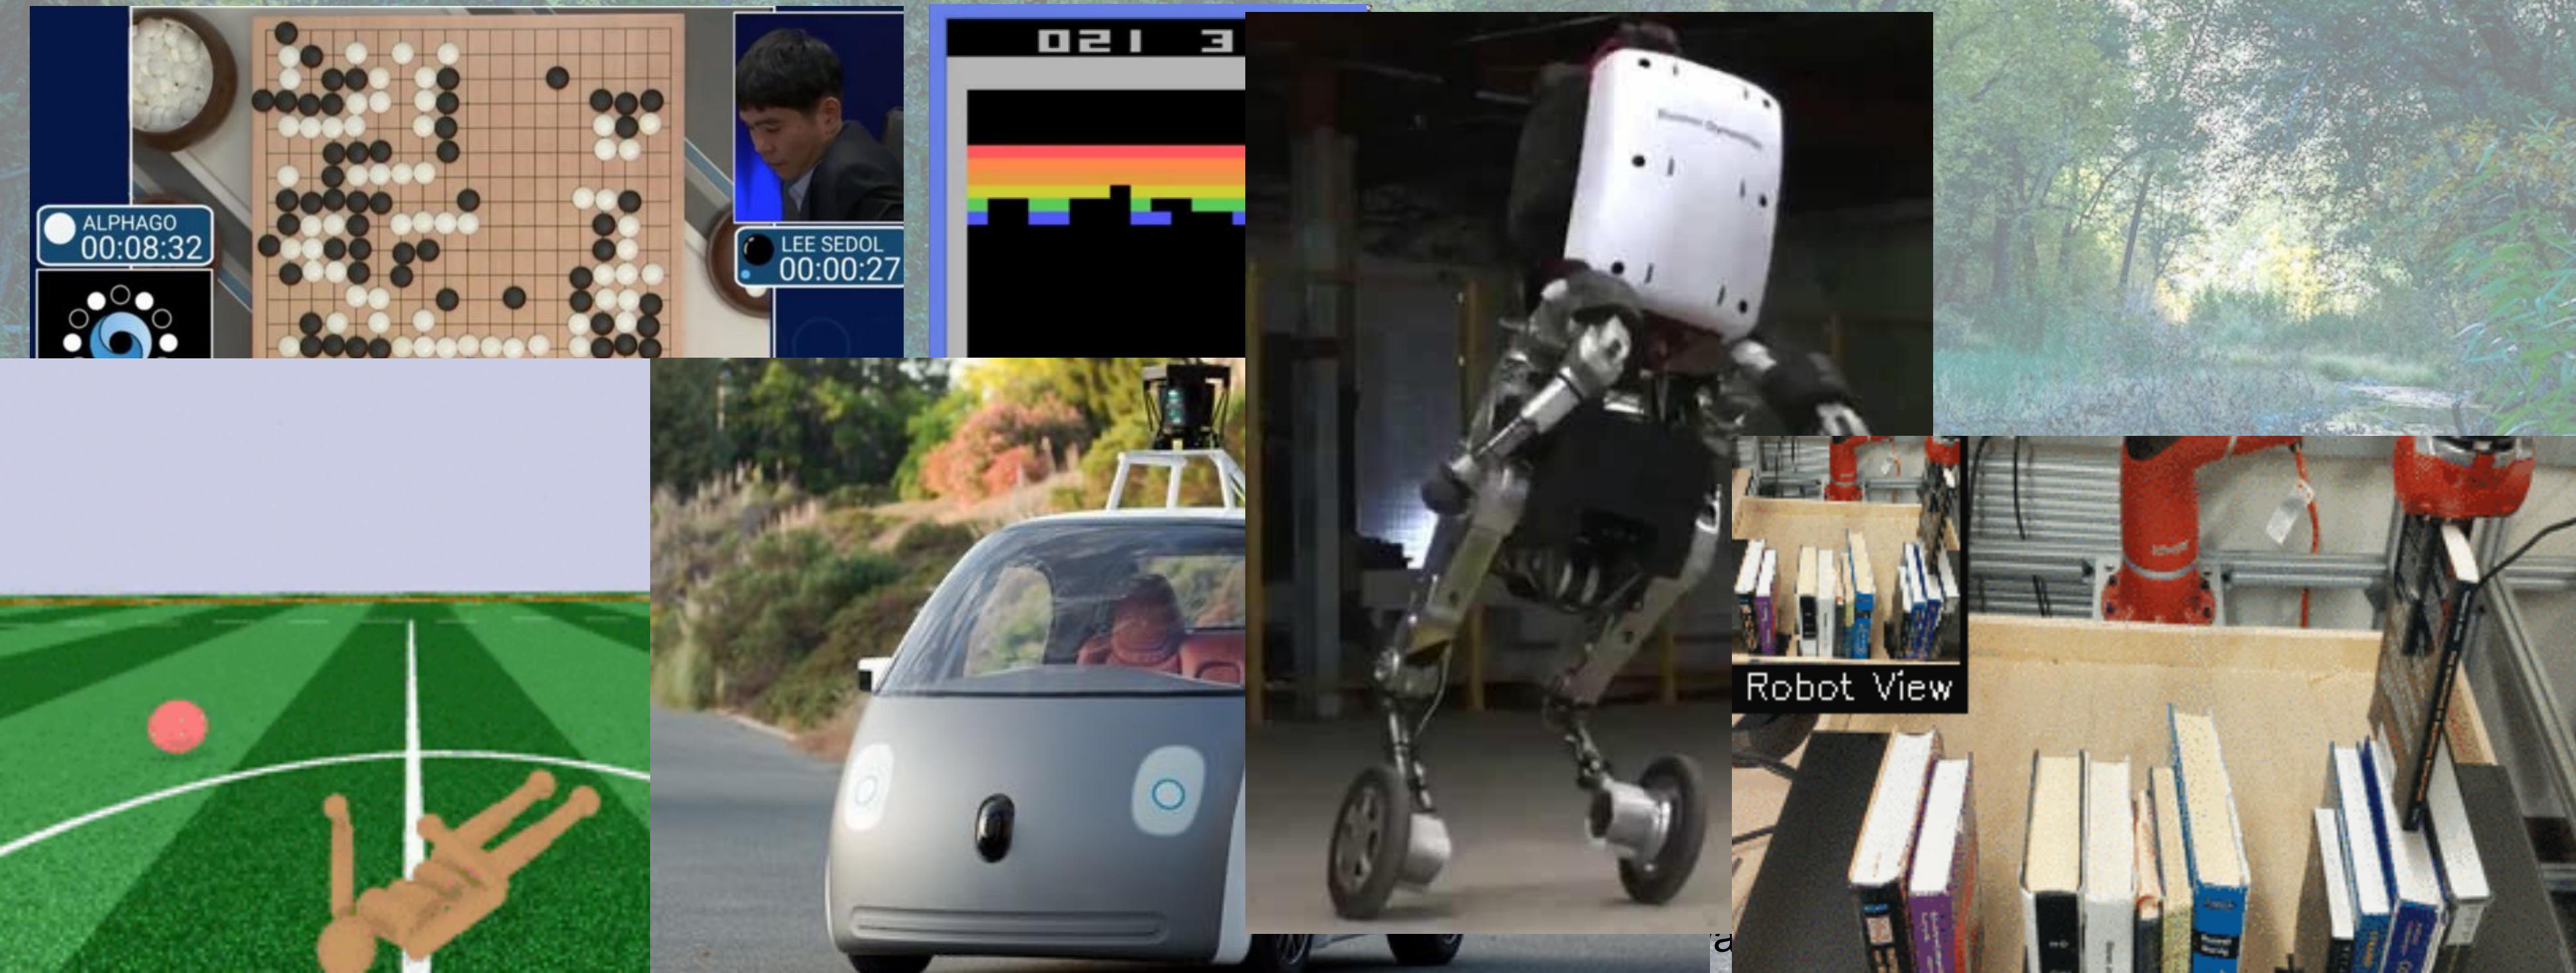
Industrial
Processes

System
Optimization

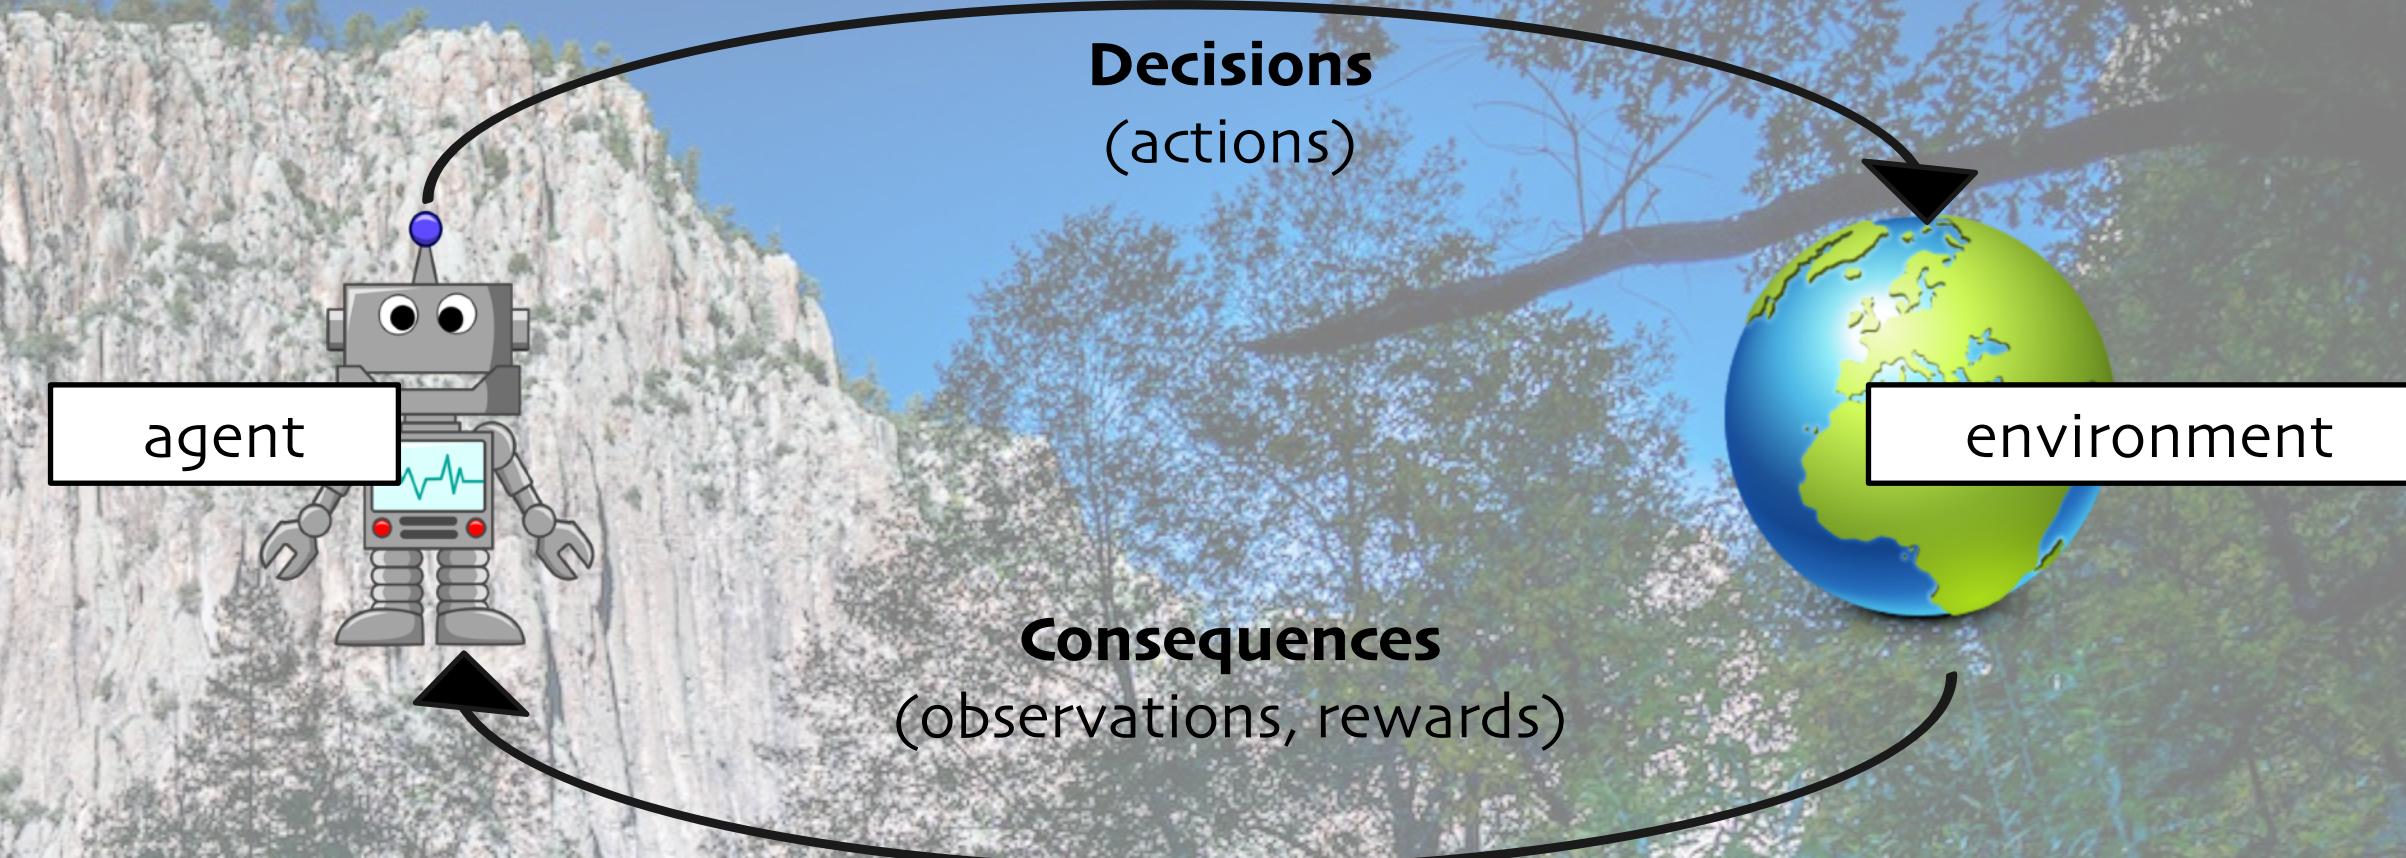
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

Industrial
Processes

System
Optimization

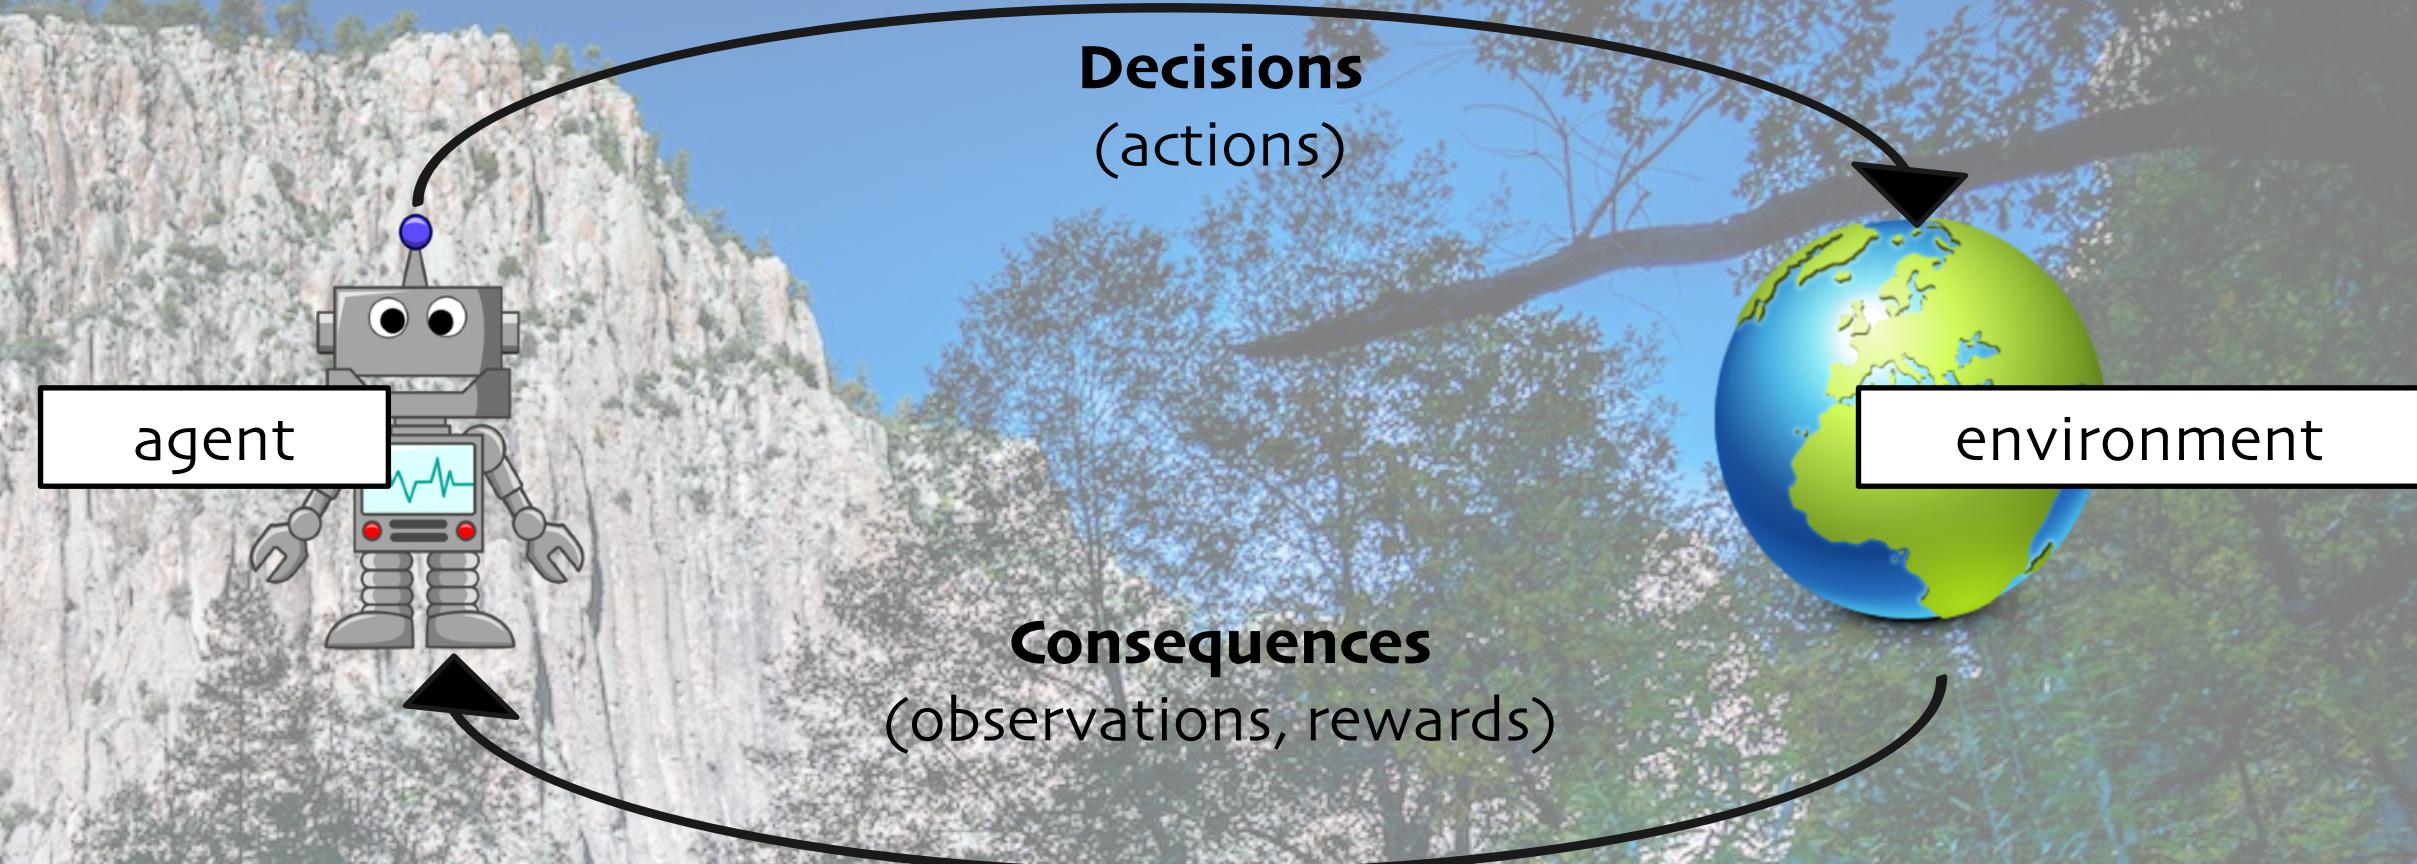
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

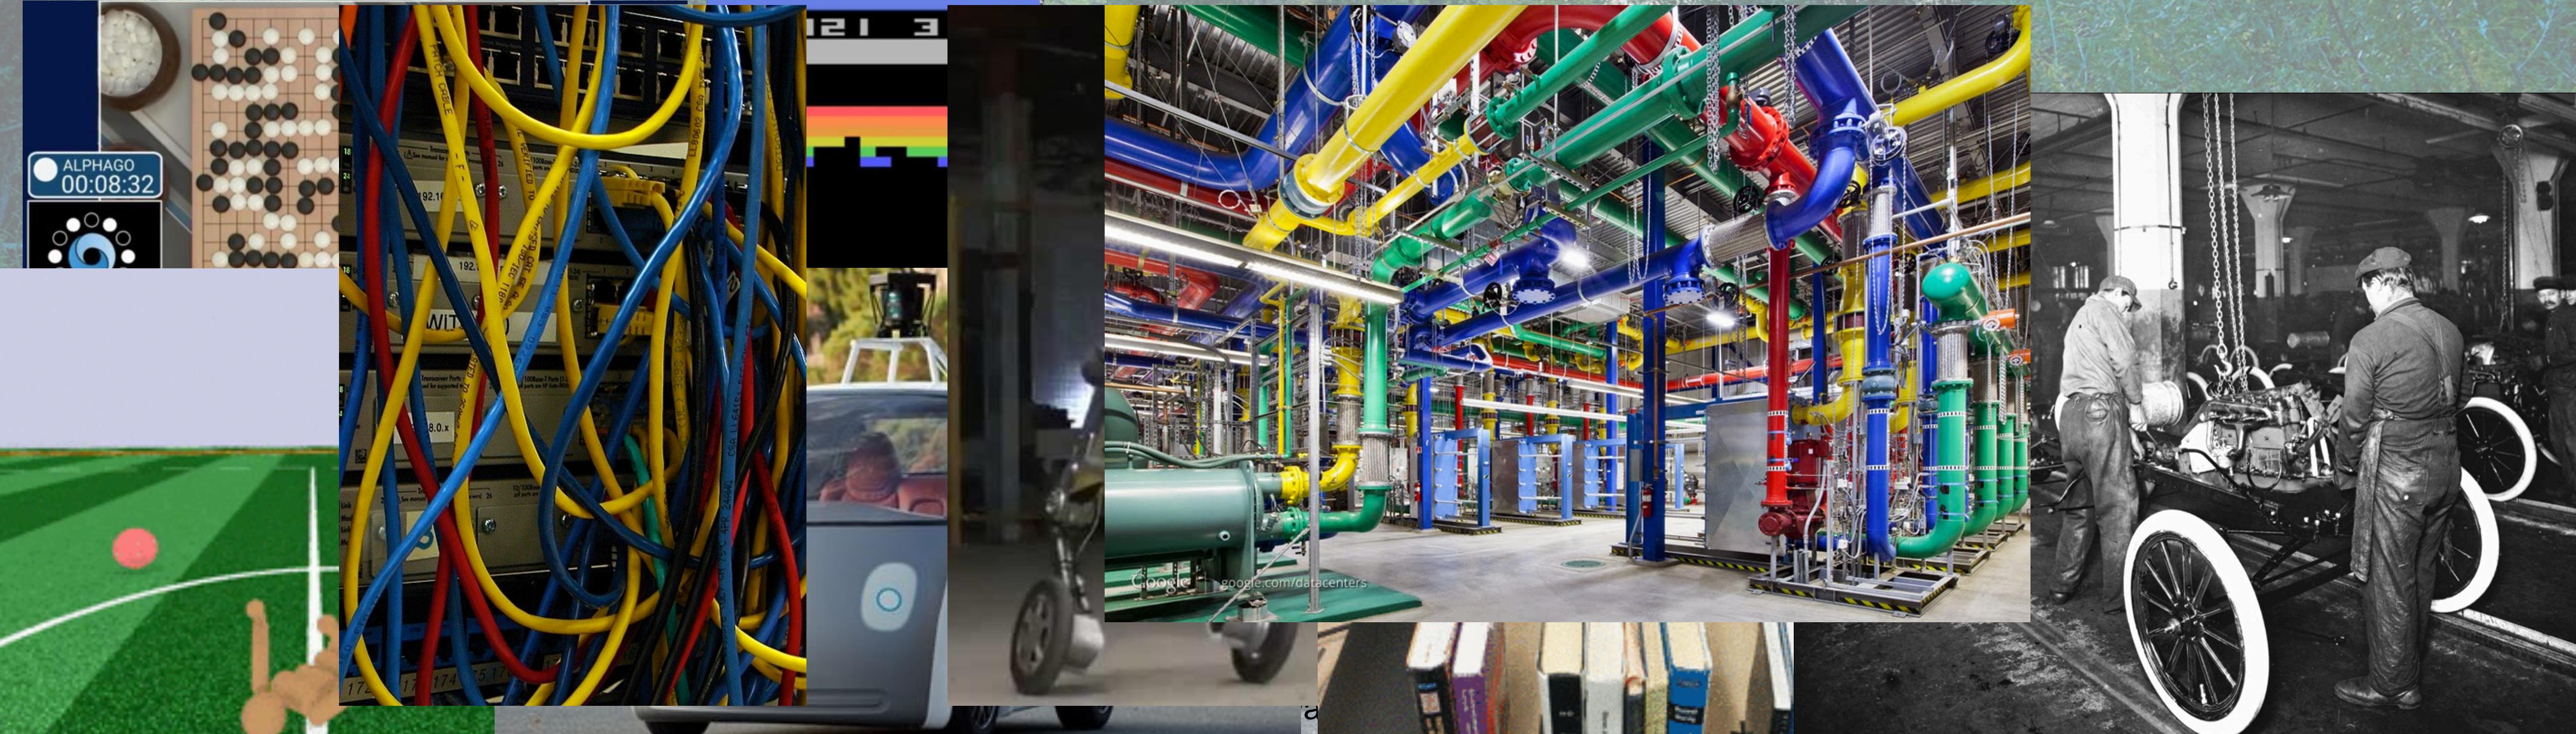
Industrial
Processes

System
Optimization

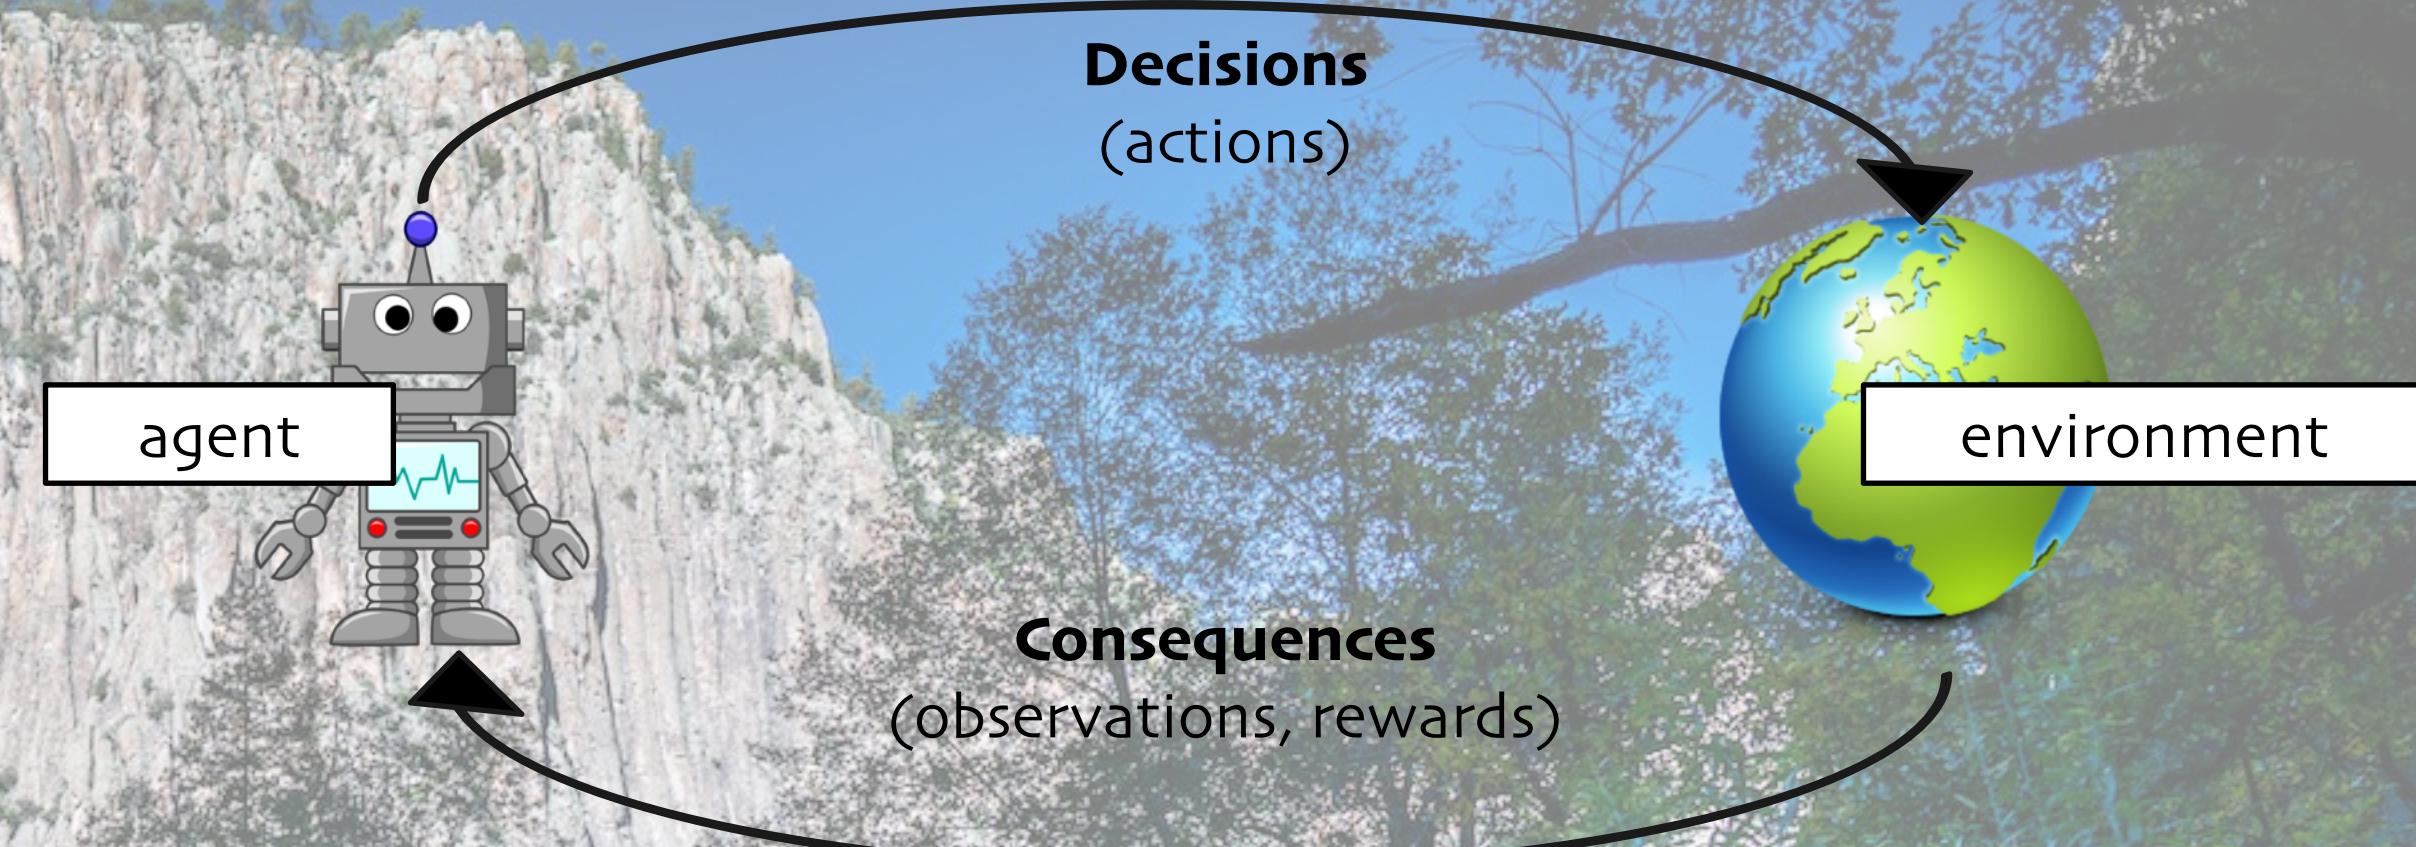
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

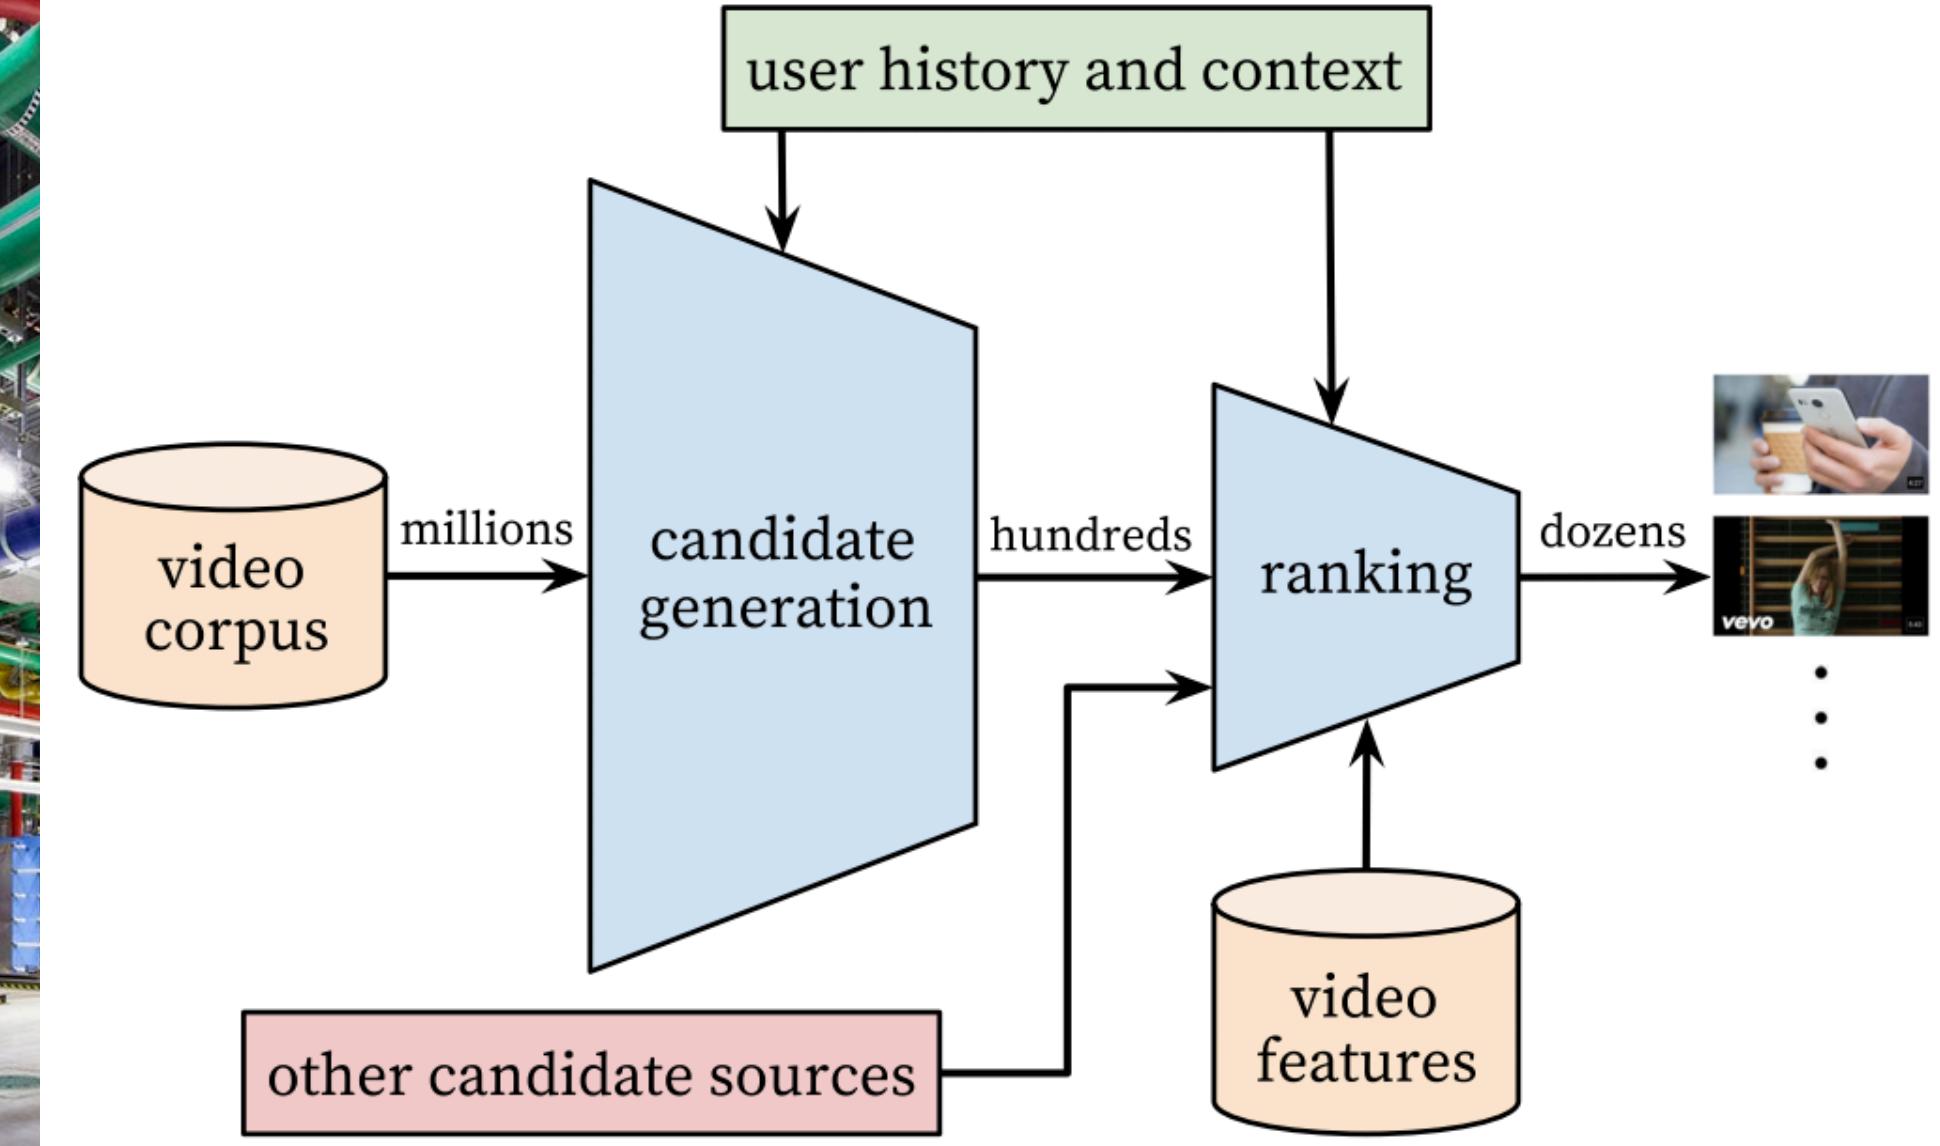
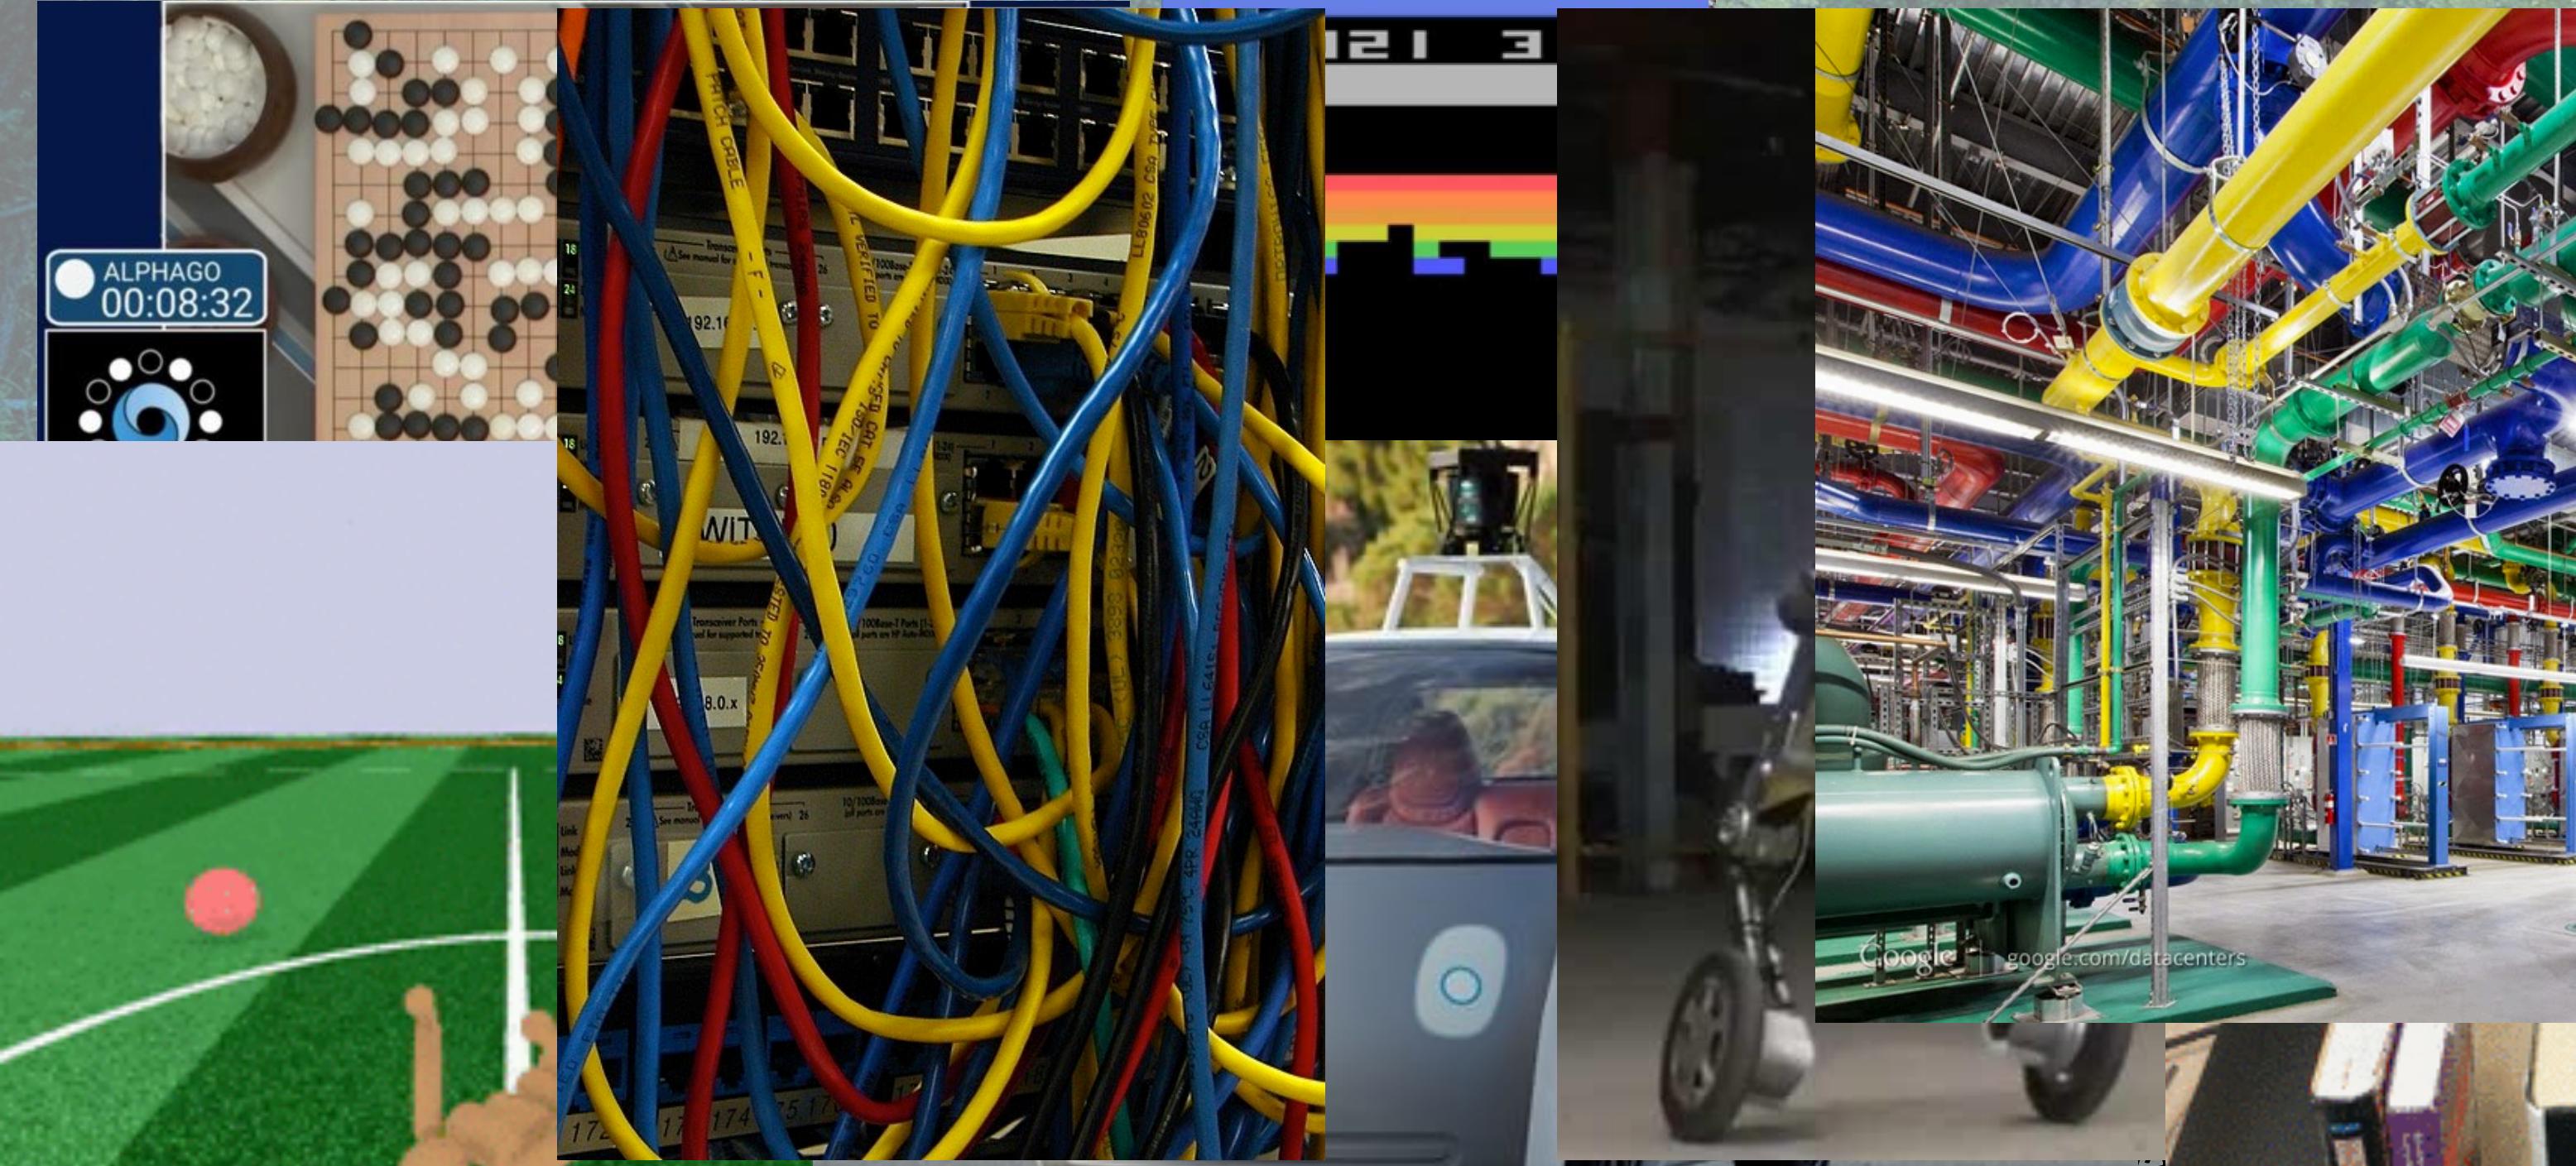
Industrial
Processes

System
Optimization

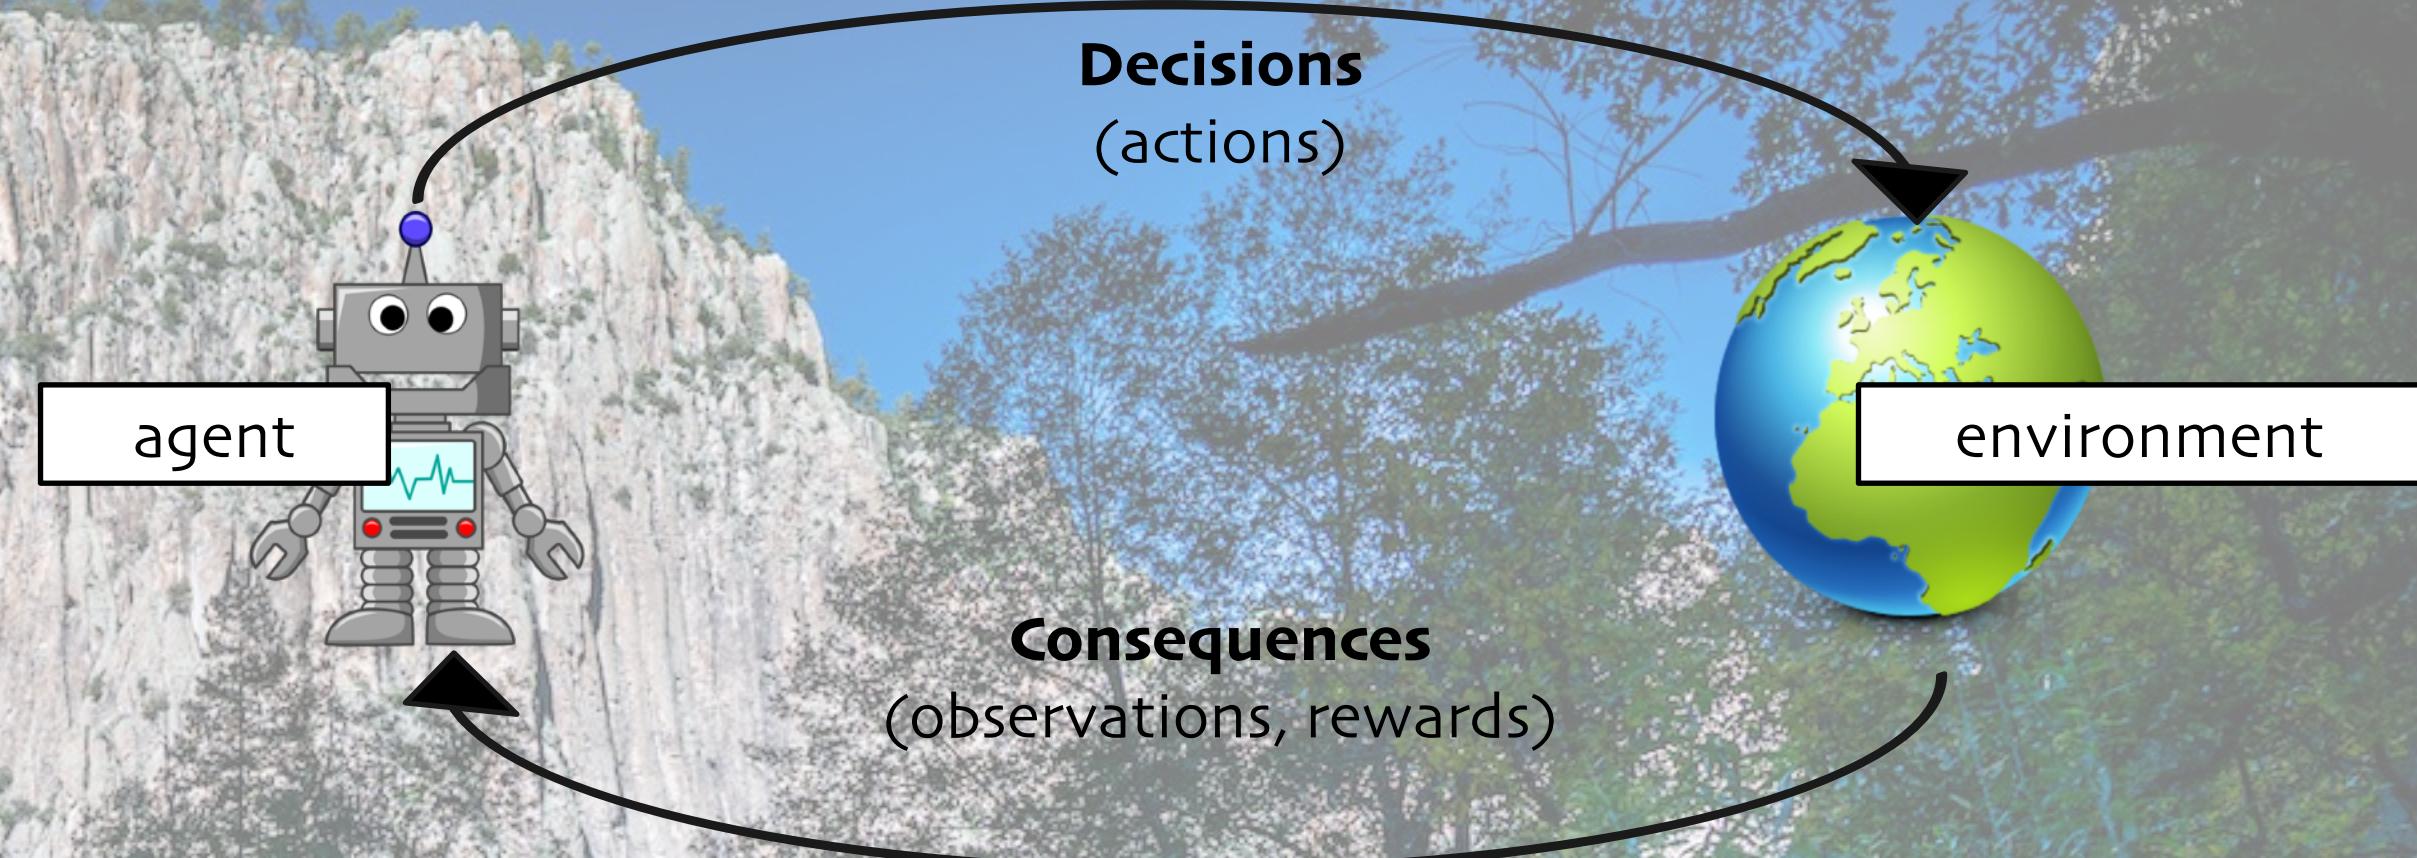
Advertising,
Recommendations

Finance

RL applications



Reinforcement Learning



Games

Robotics,
Autonomous
Vehicles

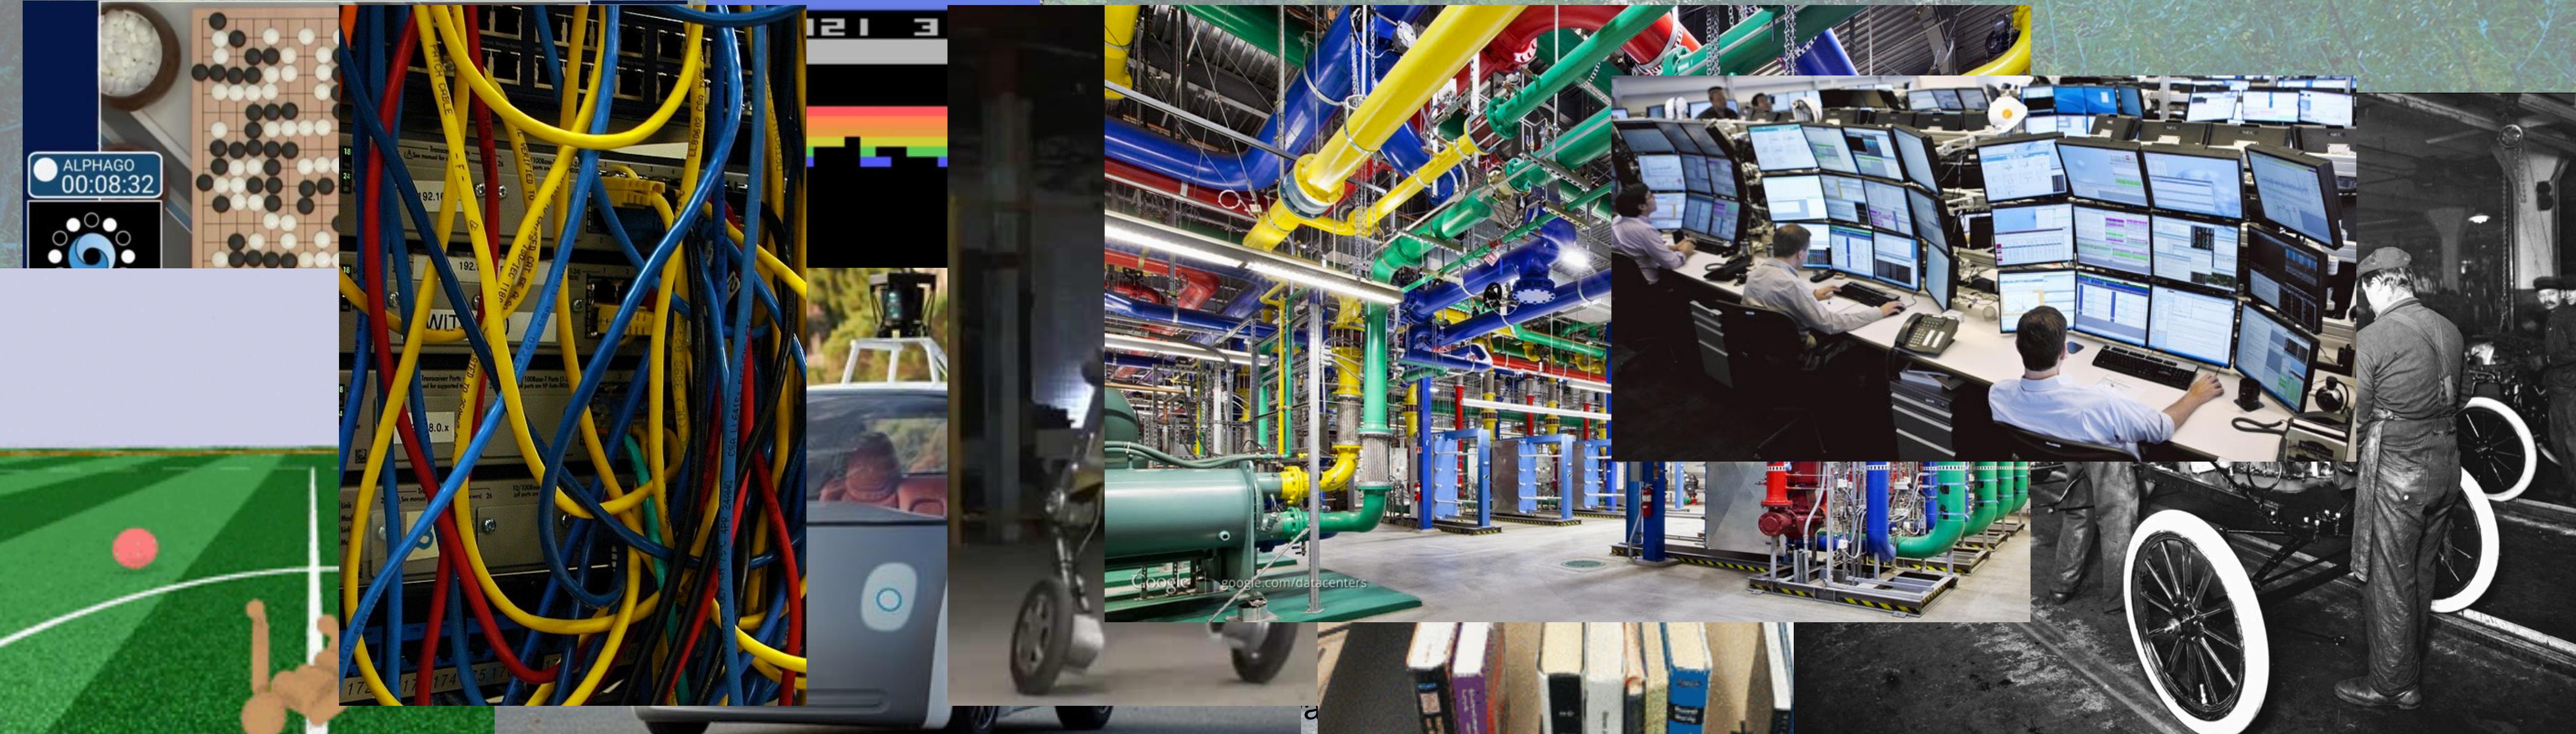
Industrial
Processes

System
Optimization

Advertising,
Recommendations

Finance

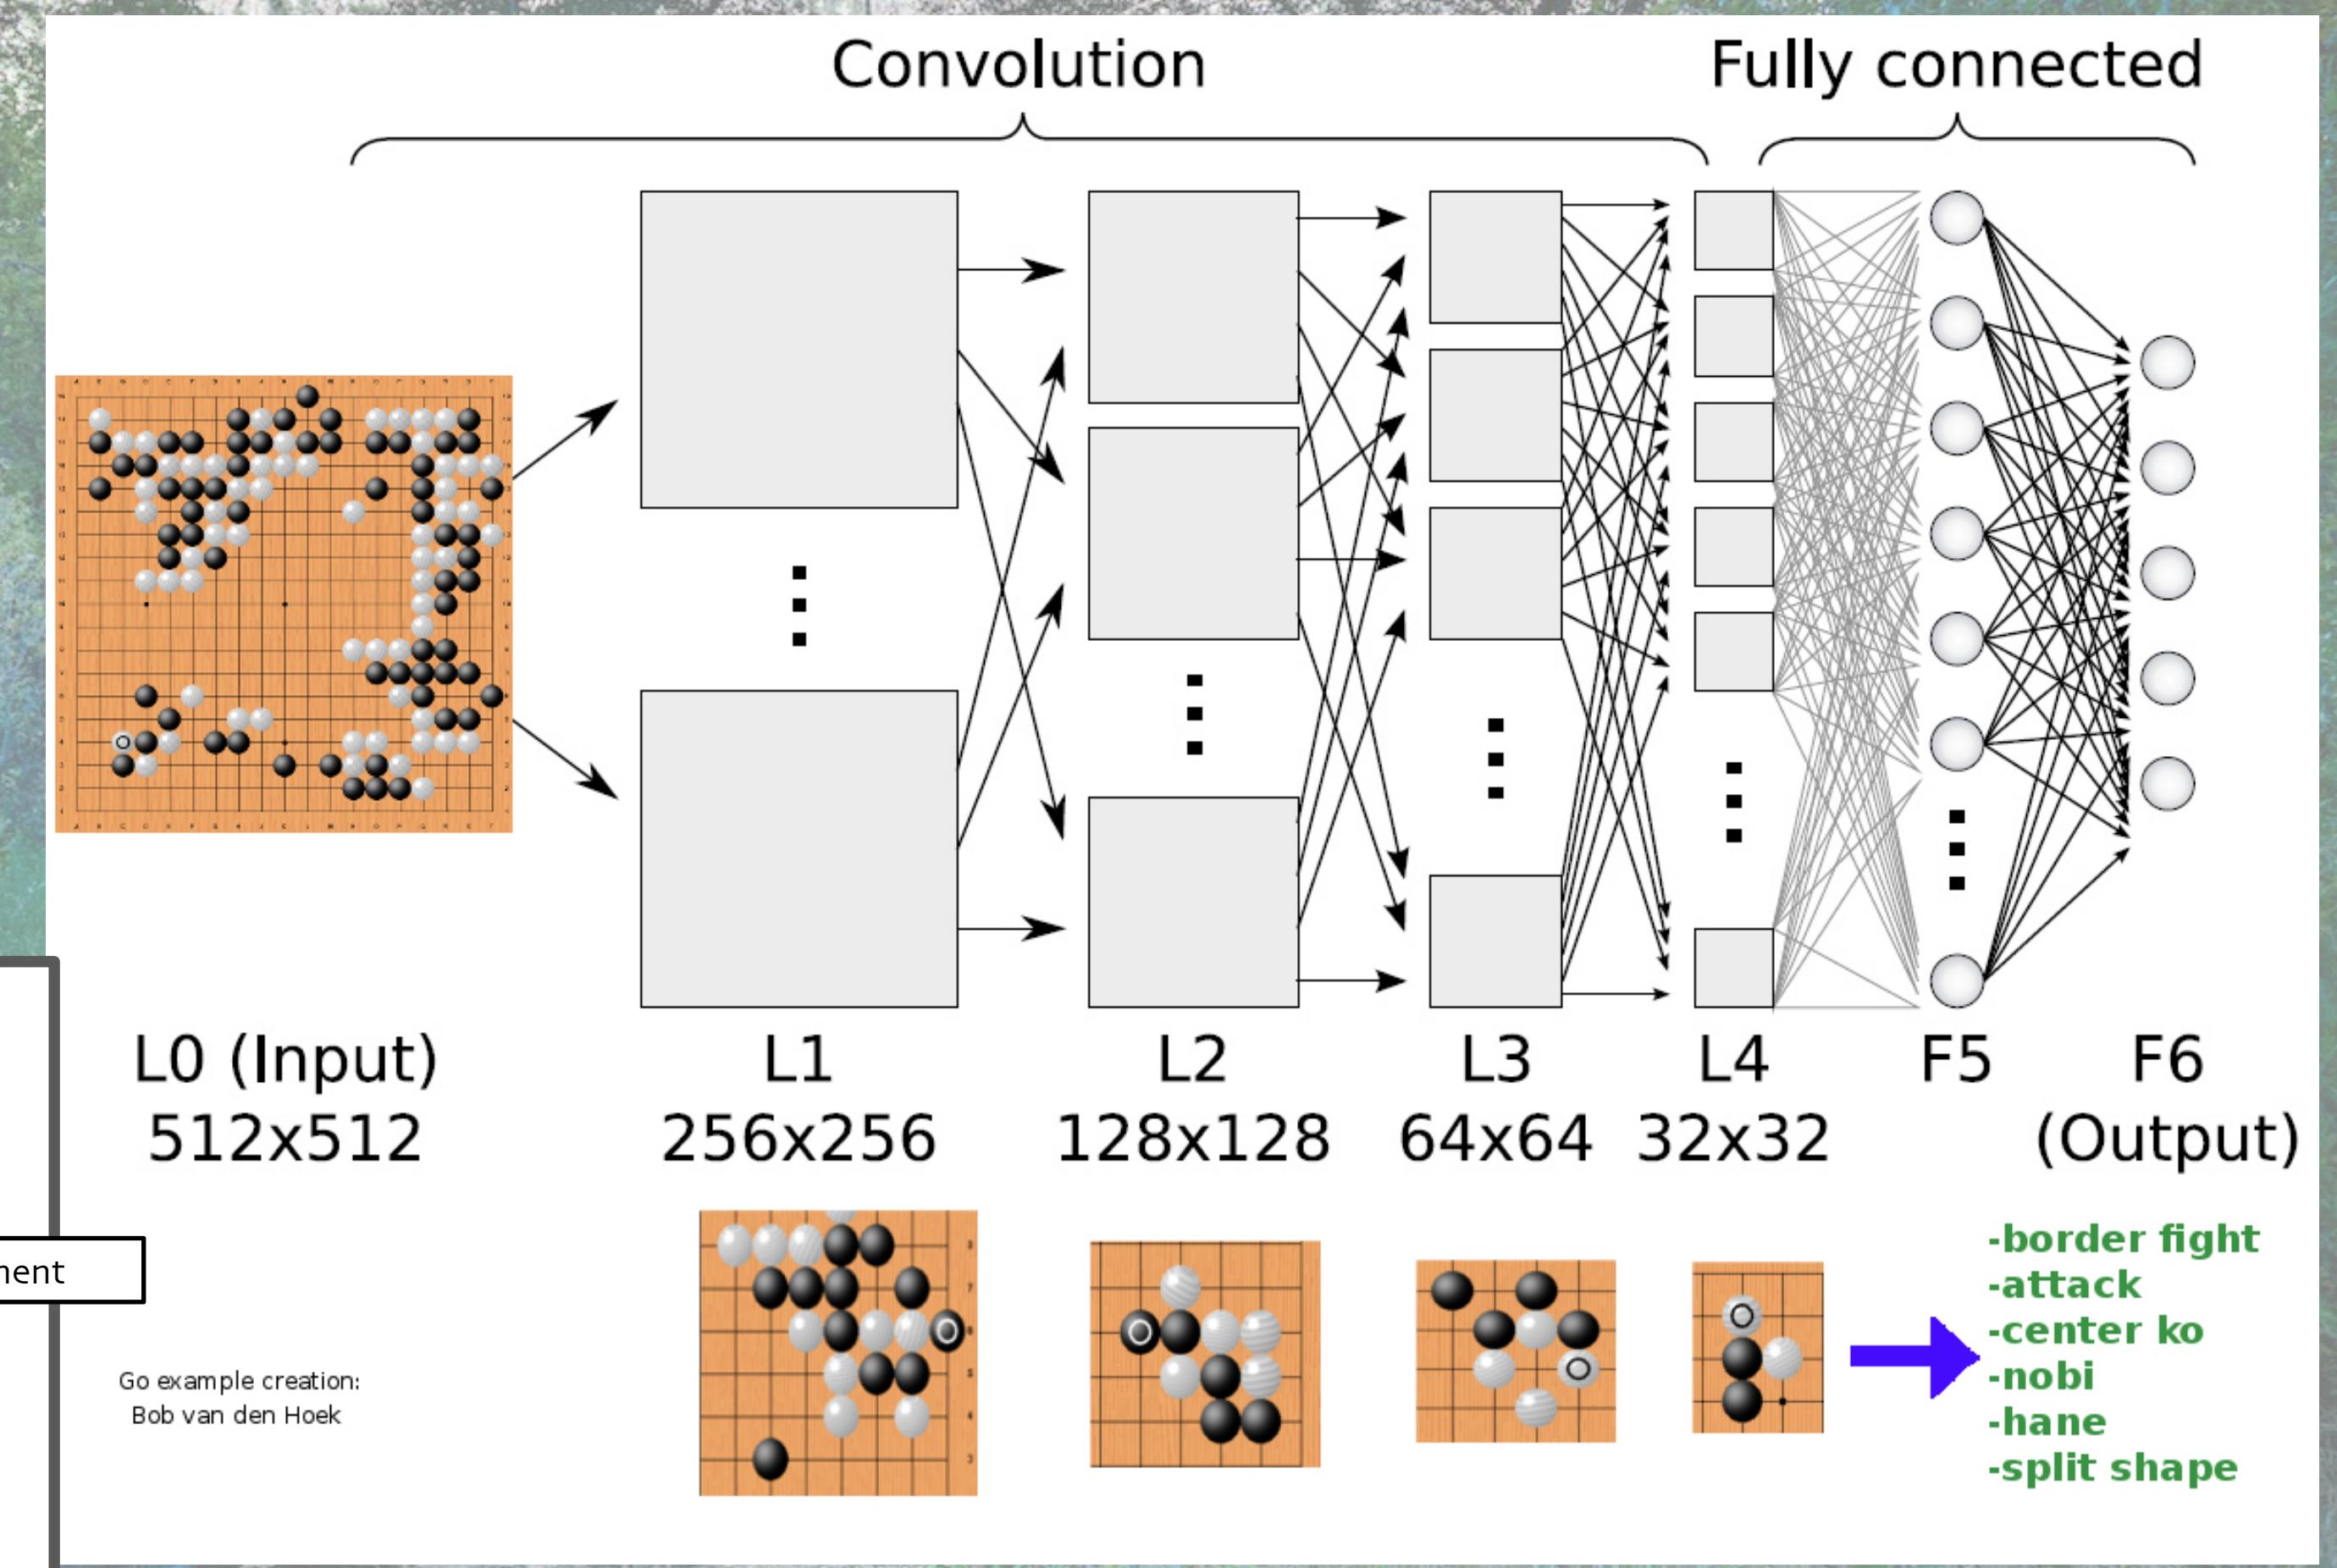
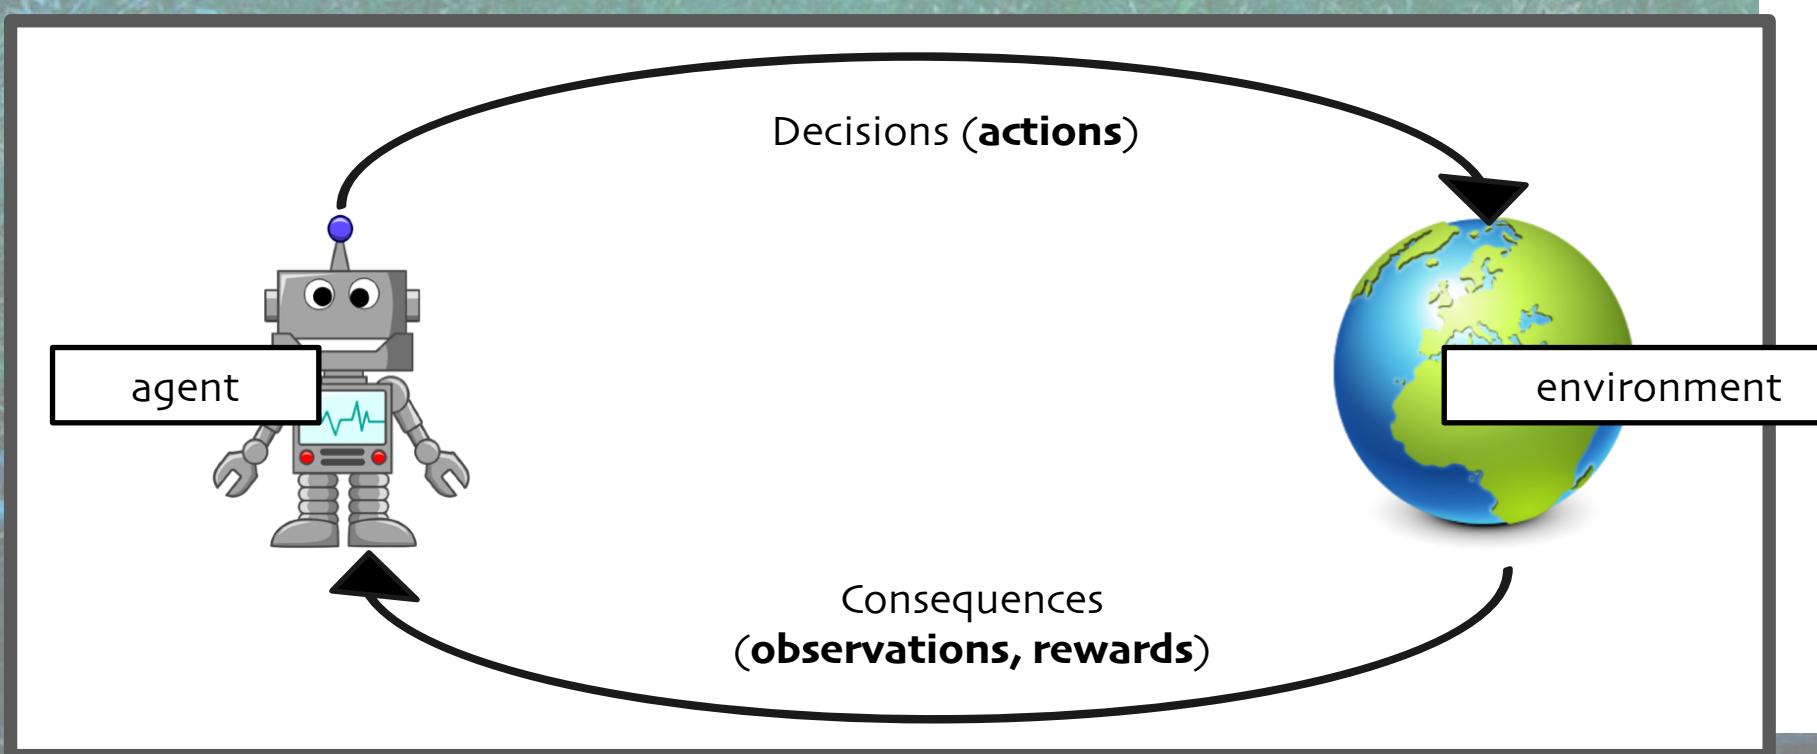
RL applications



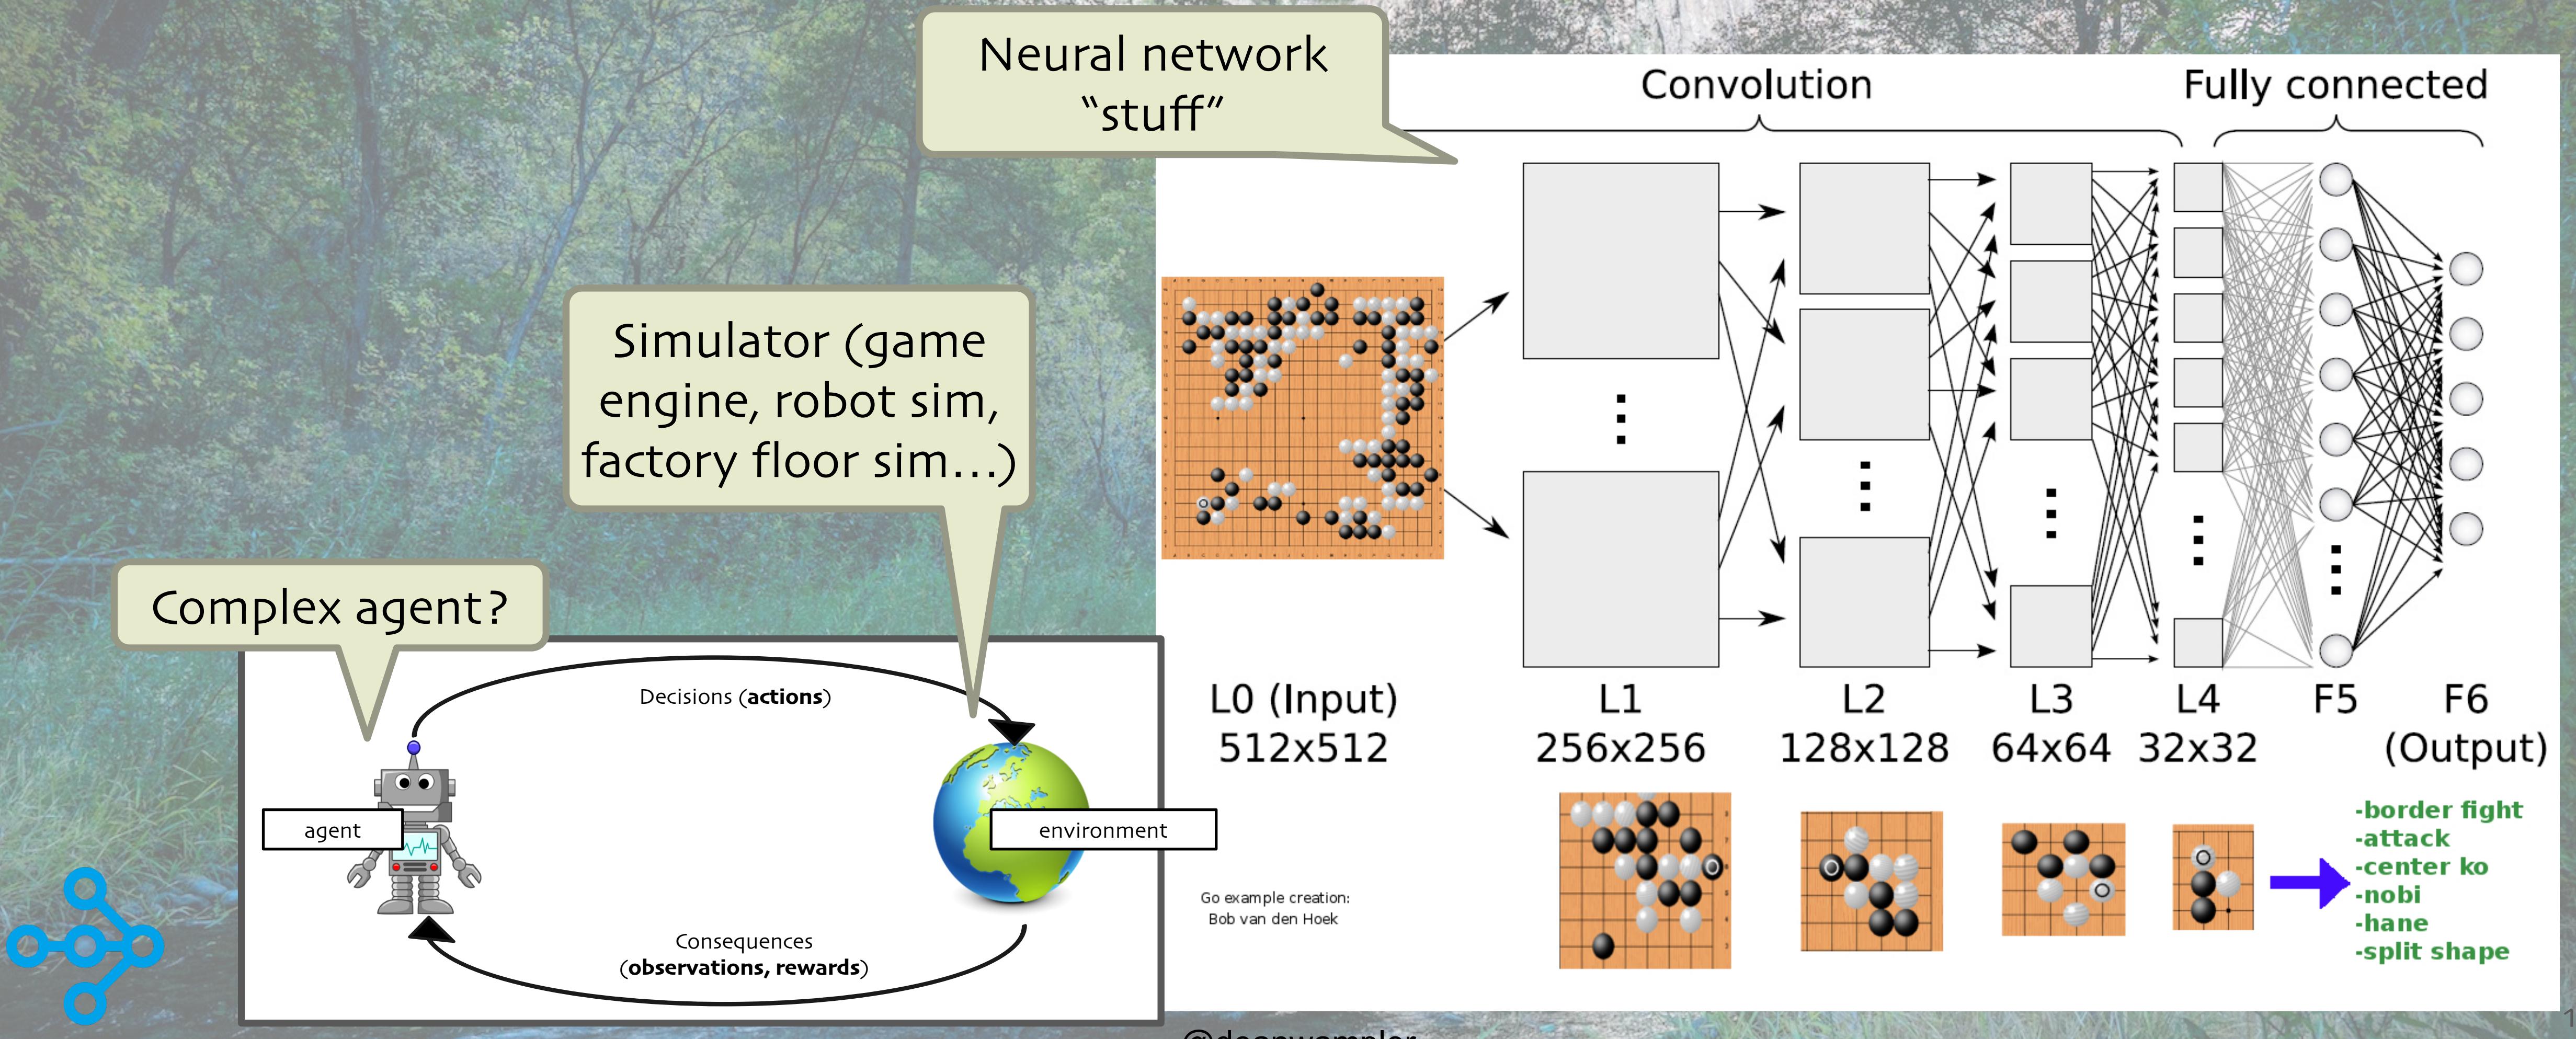
Go as a Reinforcement Learning Problem

AlphaGo (Silver et al. 2016)

- **Observations:**
 - board state
- **Actions:**
 - where to place the stones
- **Rewards:**
 - 1 if win
 - 0 otherwise



"Worst Case" Diversity of Compute Requirements



"Worst Case" Diversity of Compute Requirements

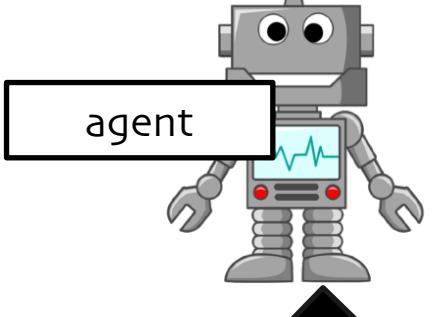
And repeated play,
over and over again,
to train for achieving
the best reward

Simulator (game
engine, robot sim,
factory floor sim...)

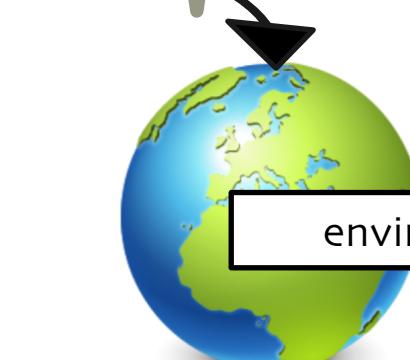
Complex agent?



Decisions (**actions**)

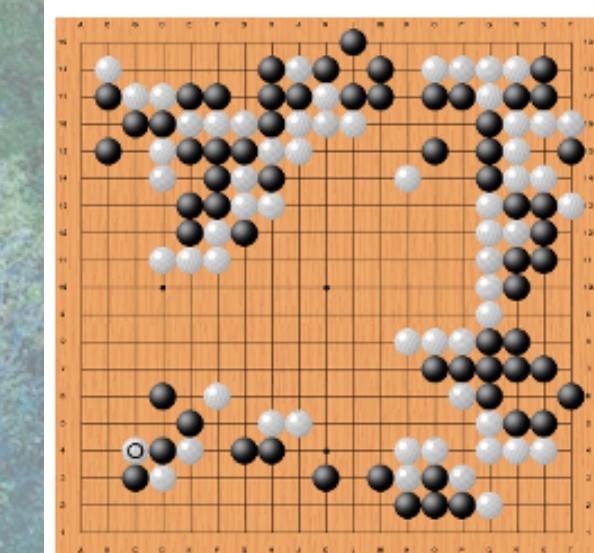


Consequences
(**observations, rewards**)

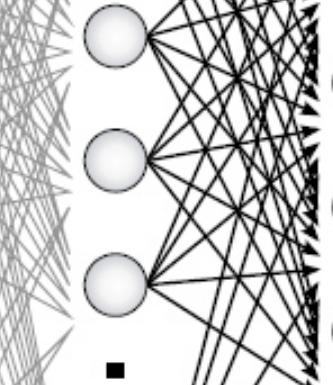
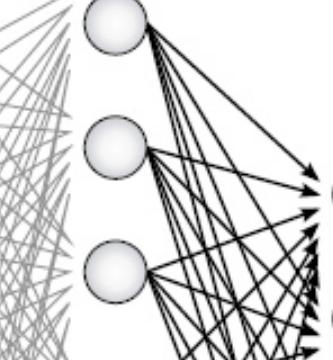
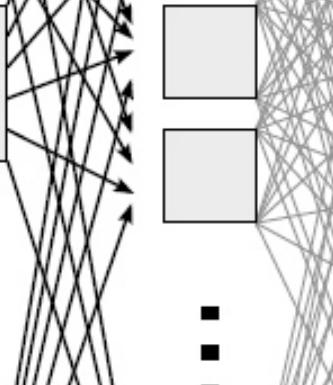
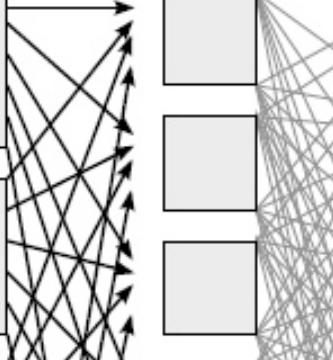
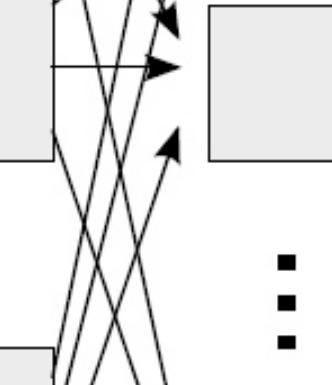
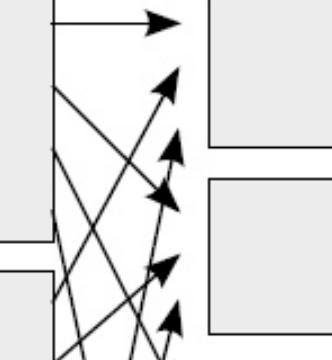
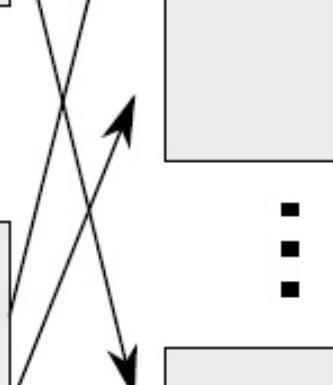
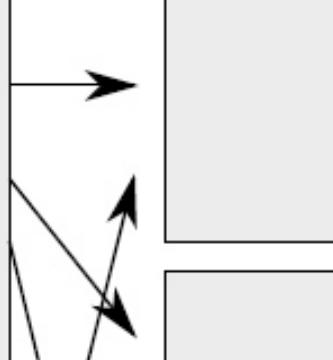
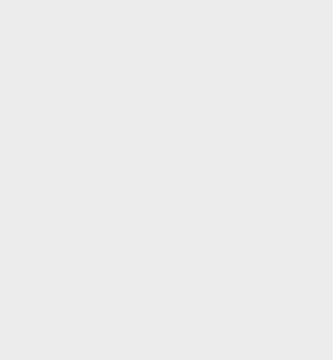


environment

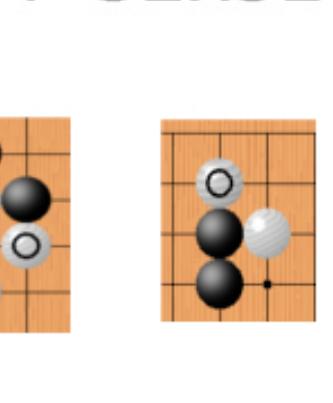
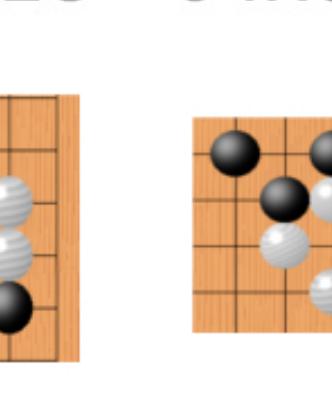
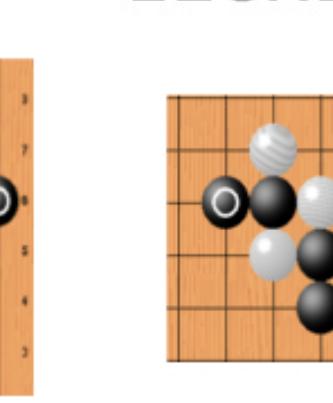
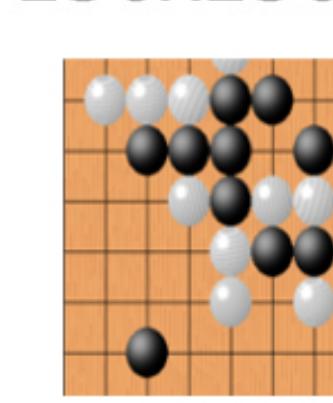
L0 (Input)
512x512



L1
256x256



L2
128x128



-border fight
-attack
-center ko
-nobi
-hane
-split shape

Convolution

Fully connected

L3
64x64

L4
32x32

L5

F6
(Output)



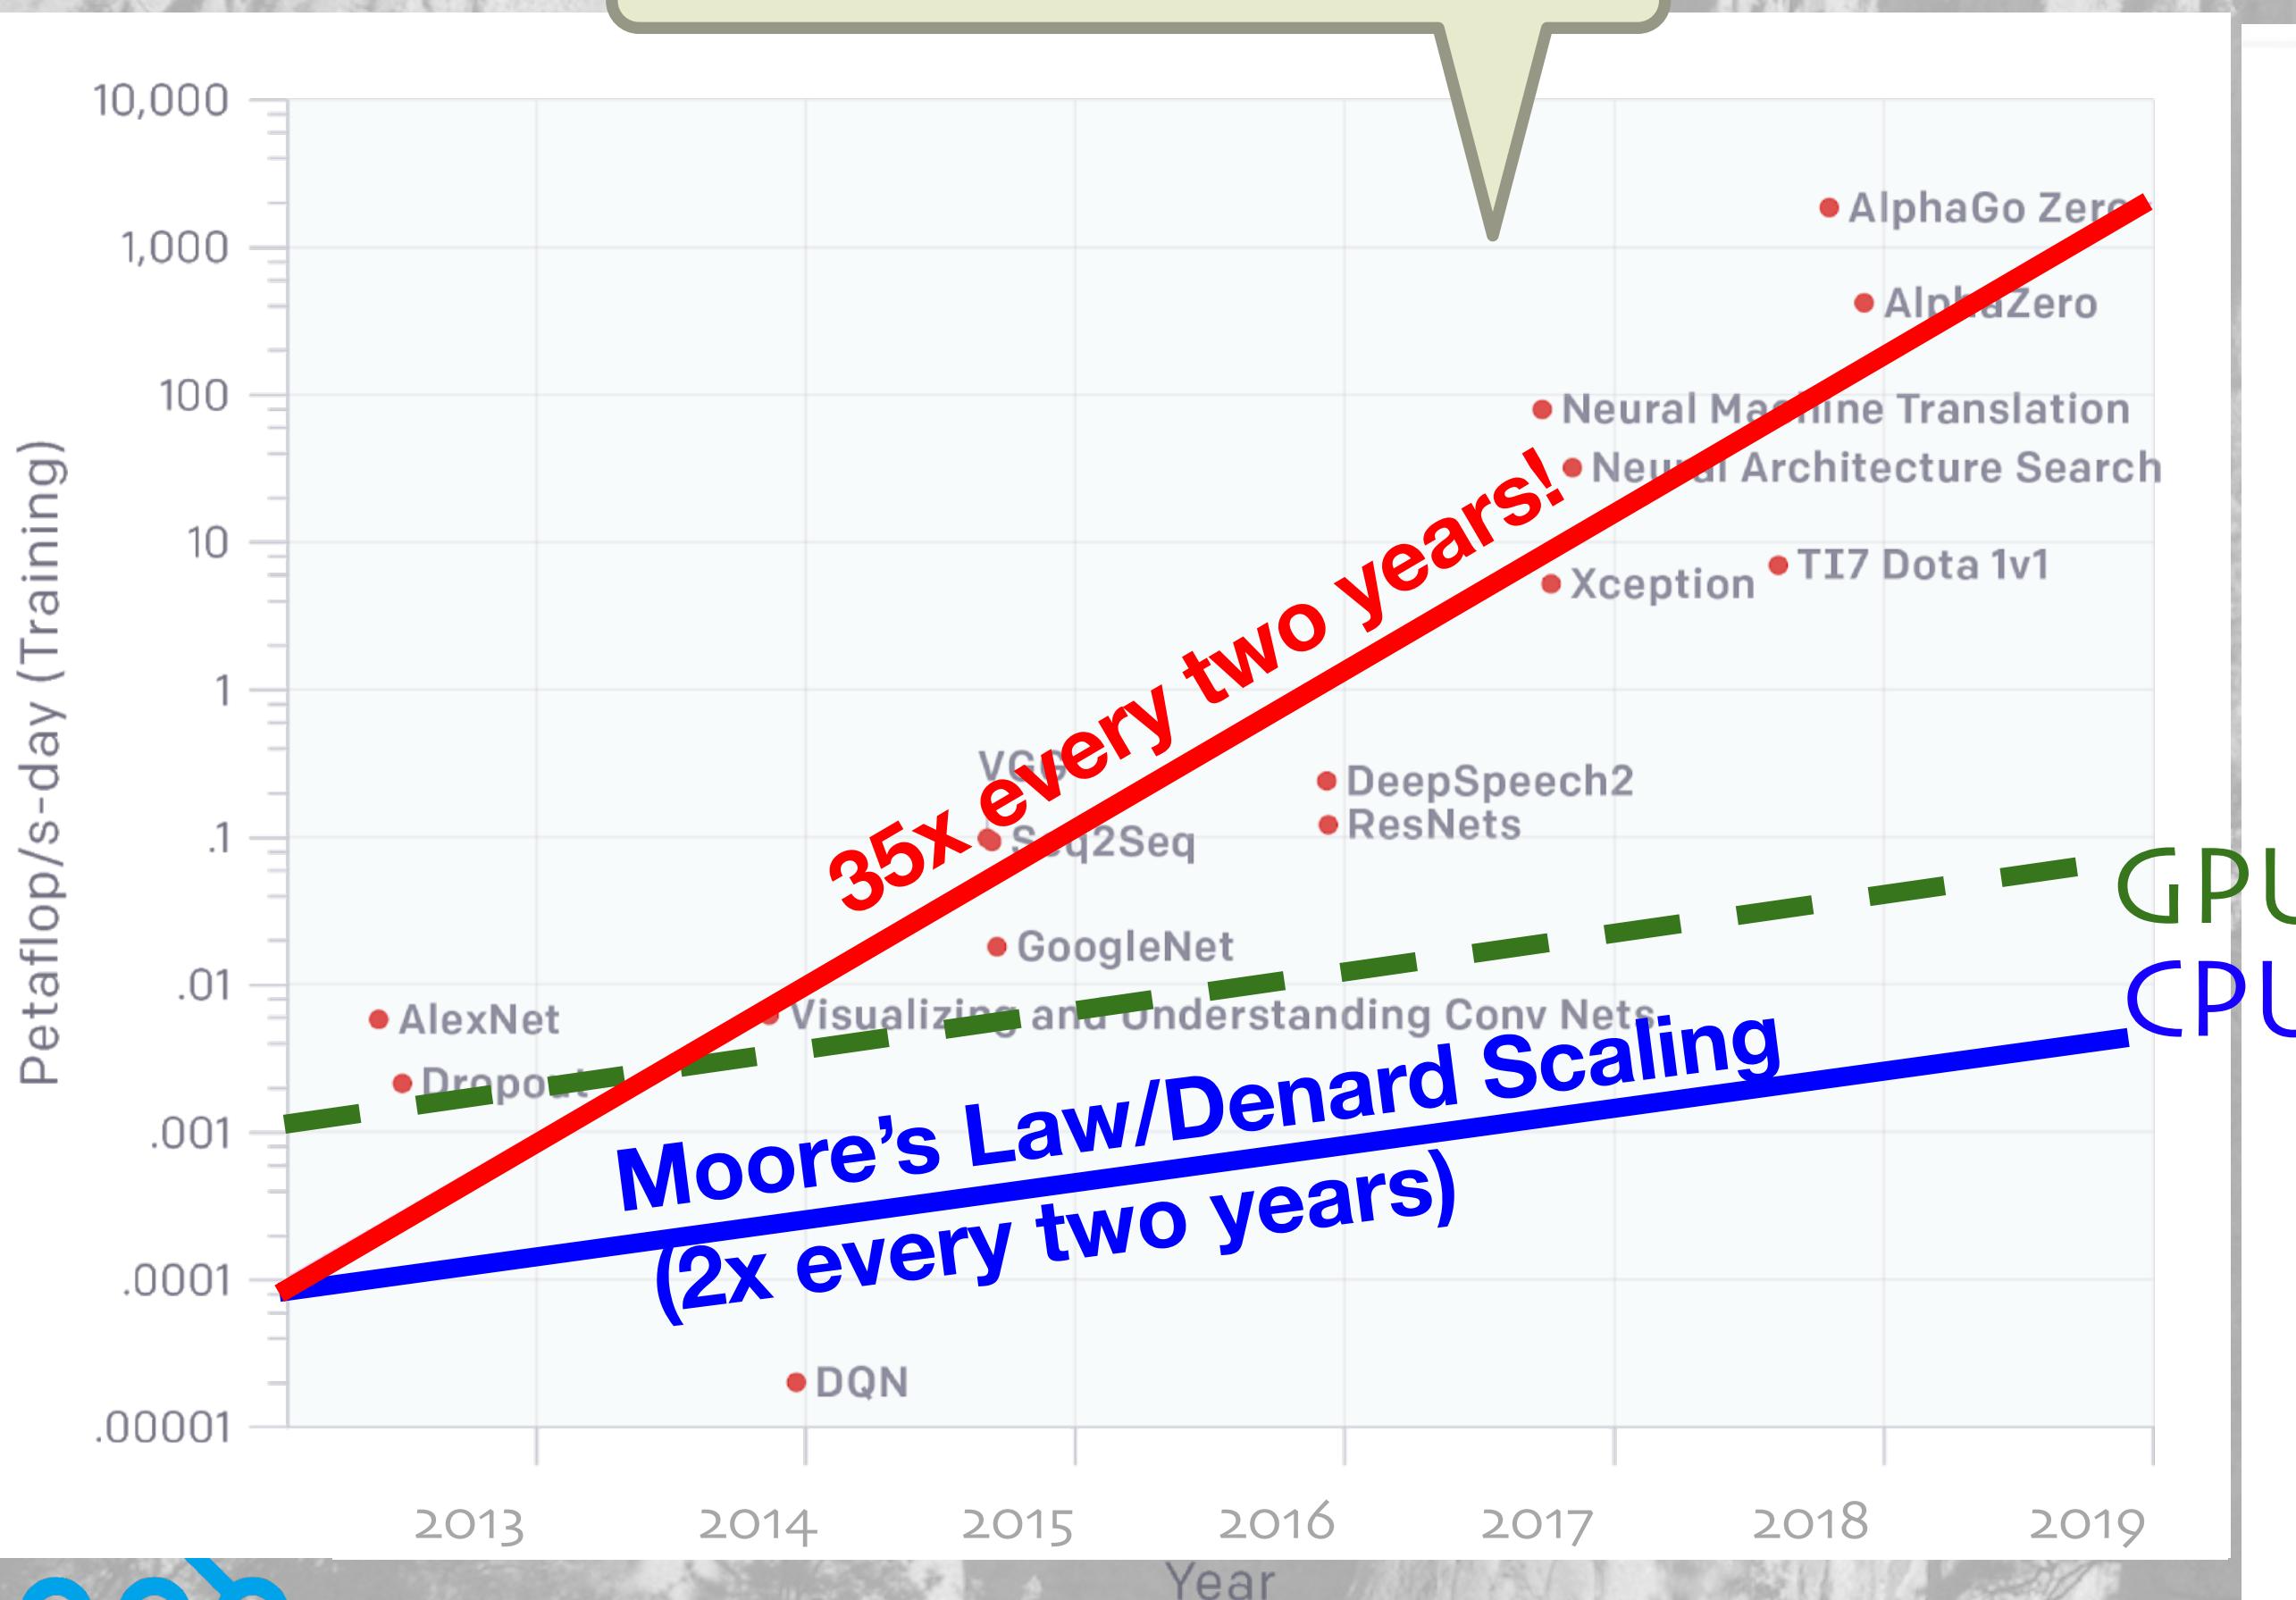
Ray to the Rescue



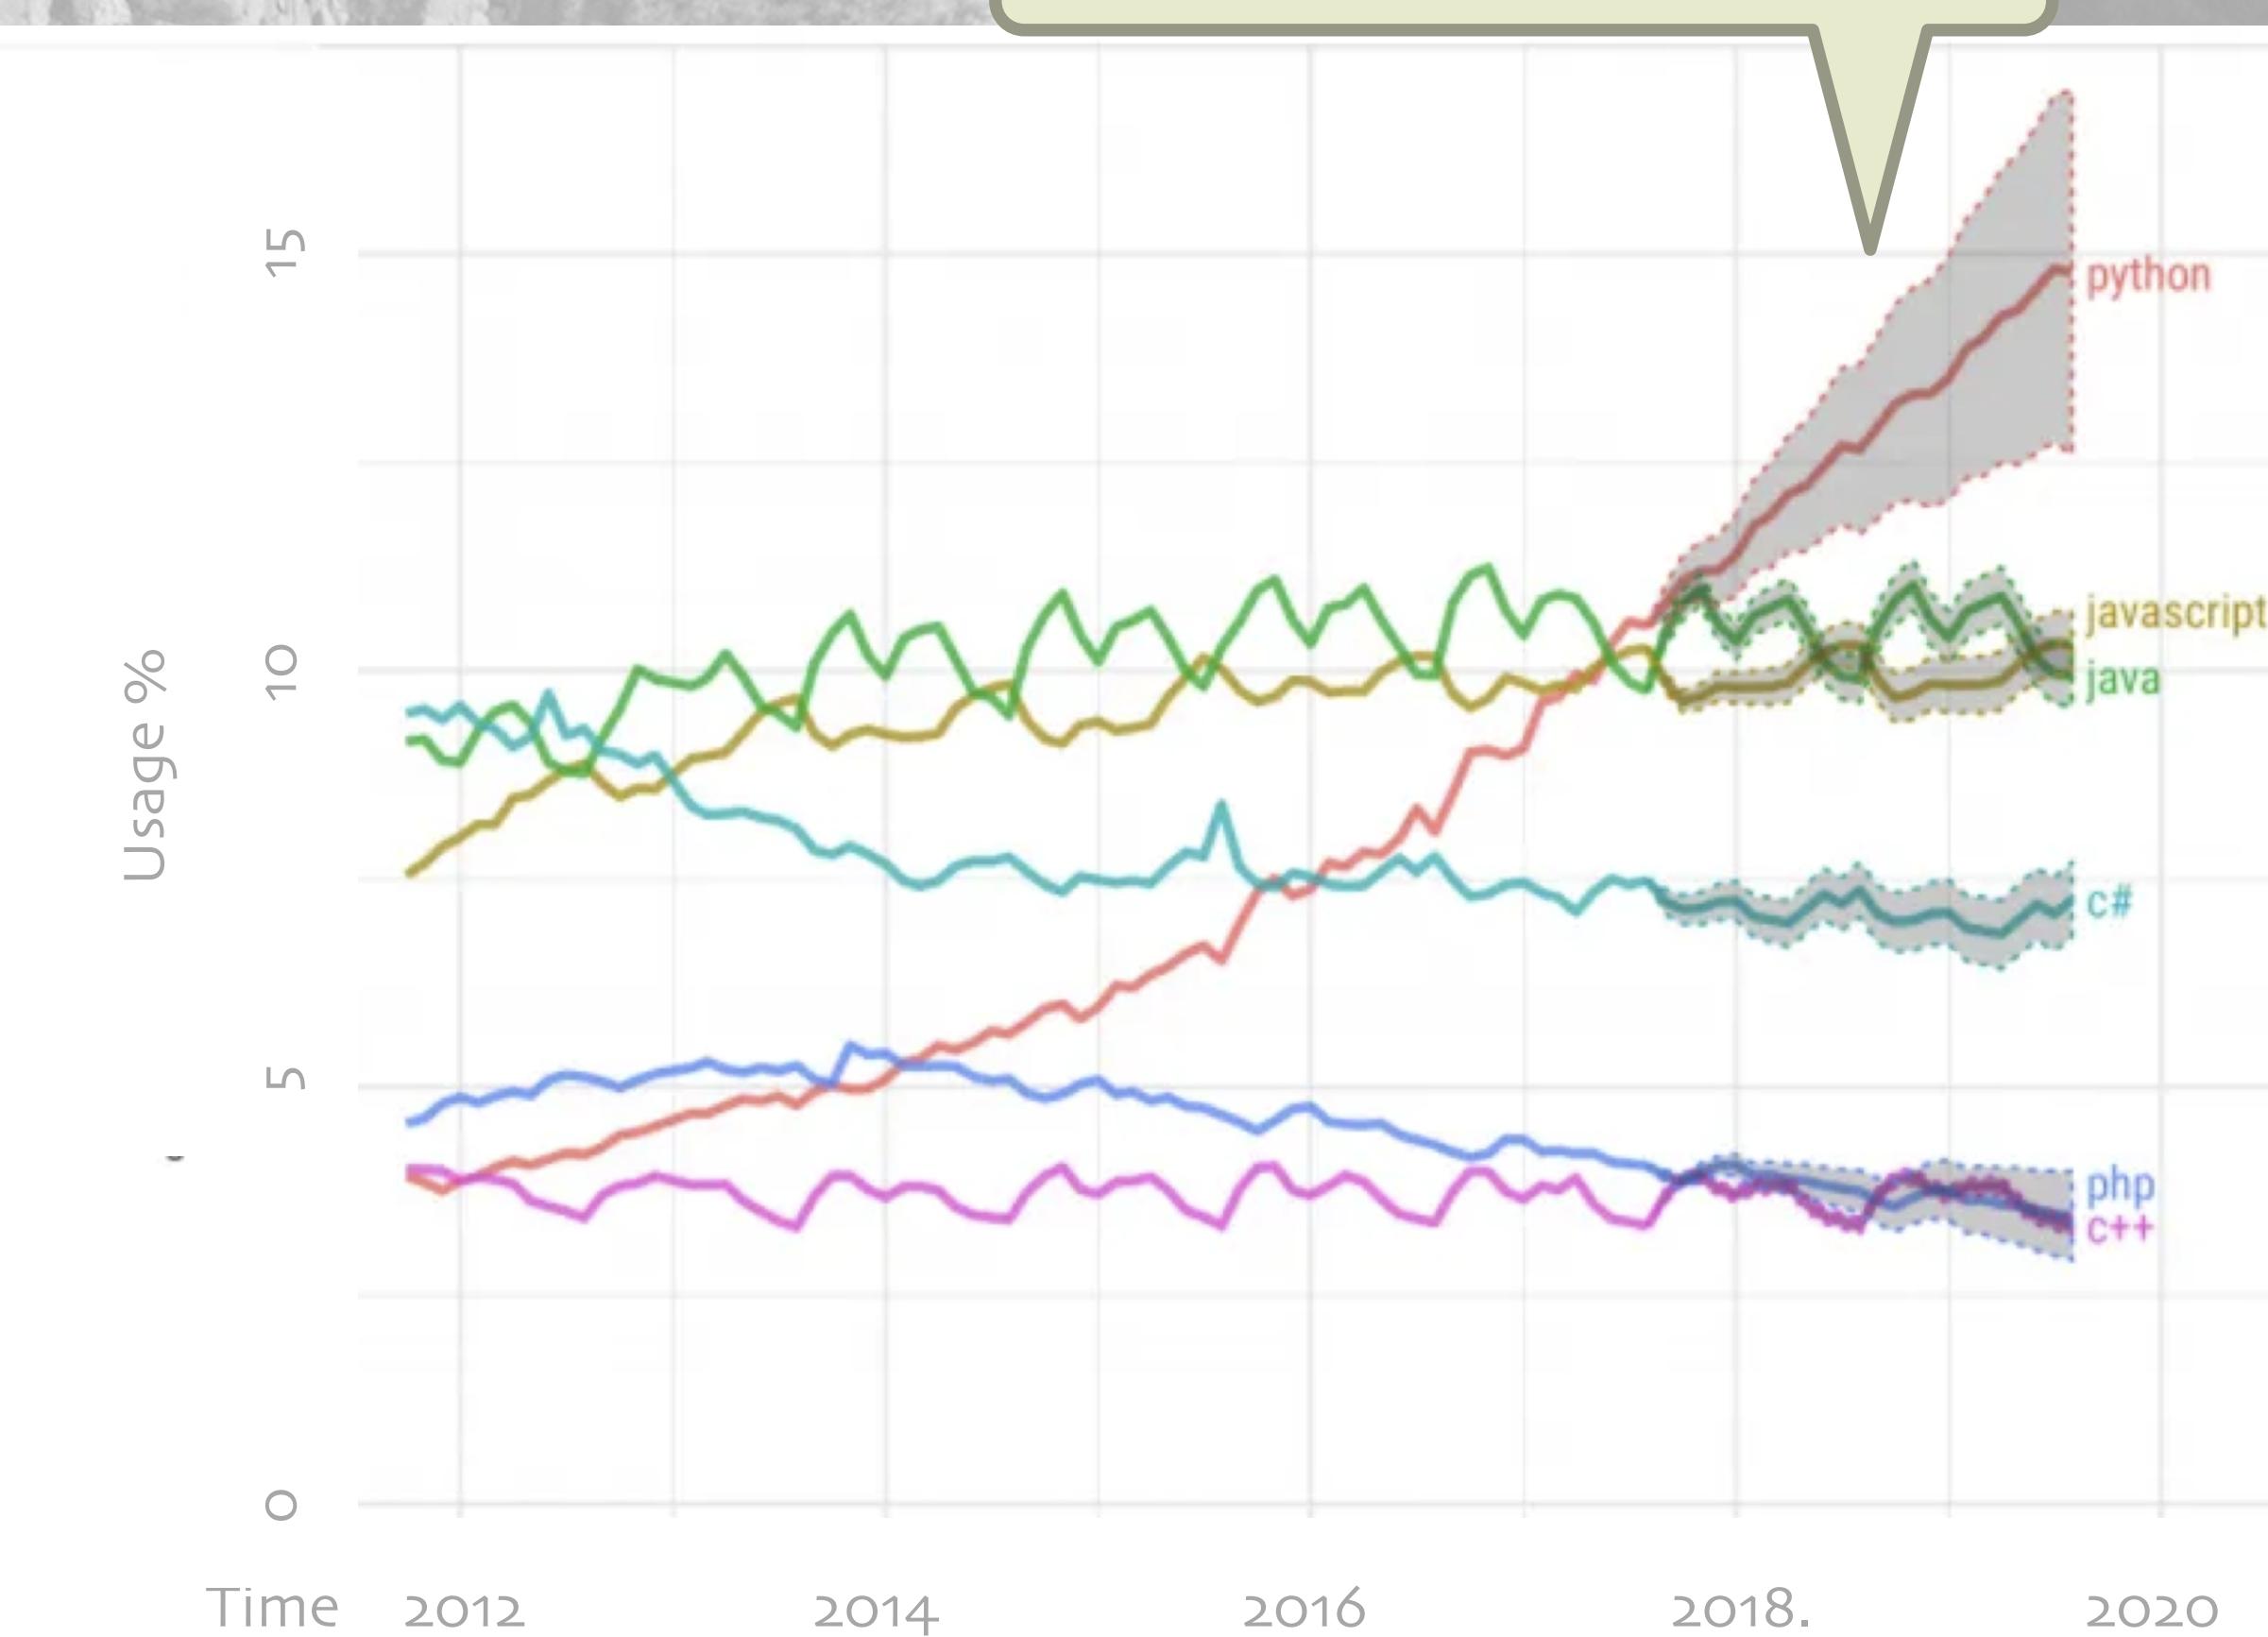
@deanwampler

Two Major Trends

Model sizes and therefore compute requirements outstripping Moore's Law



Python growth driven by ML/AI and other data science workloads

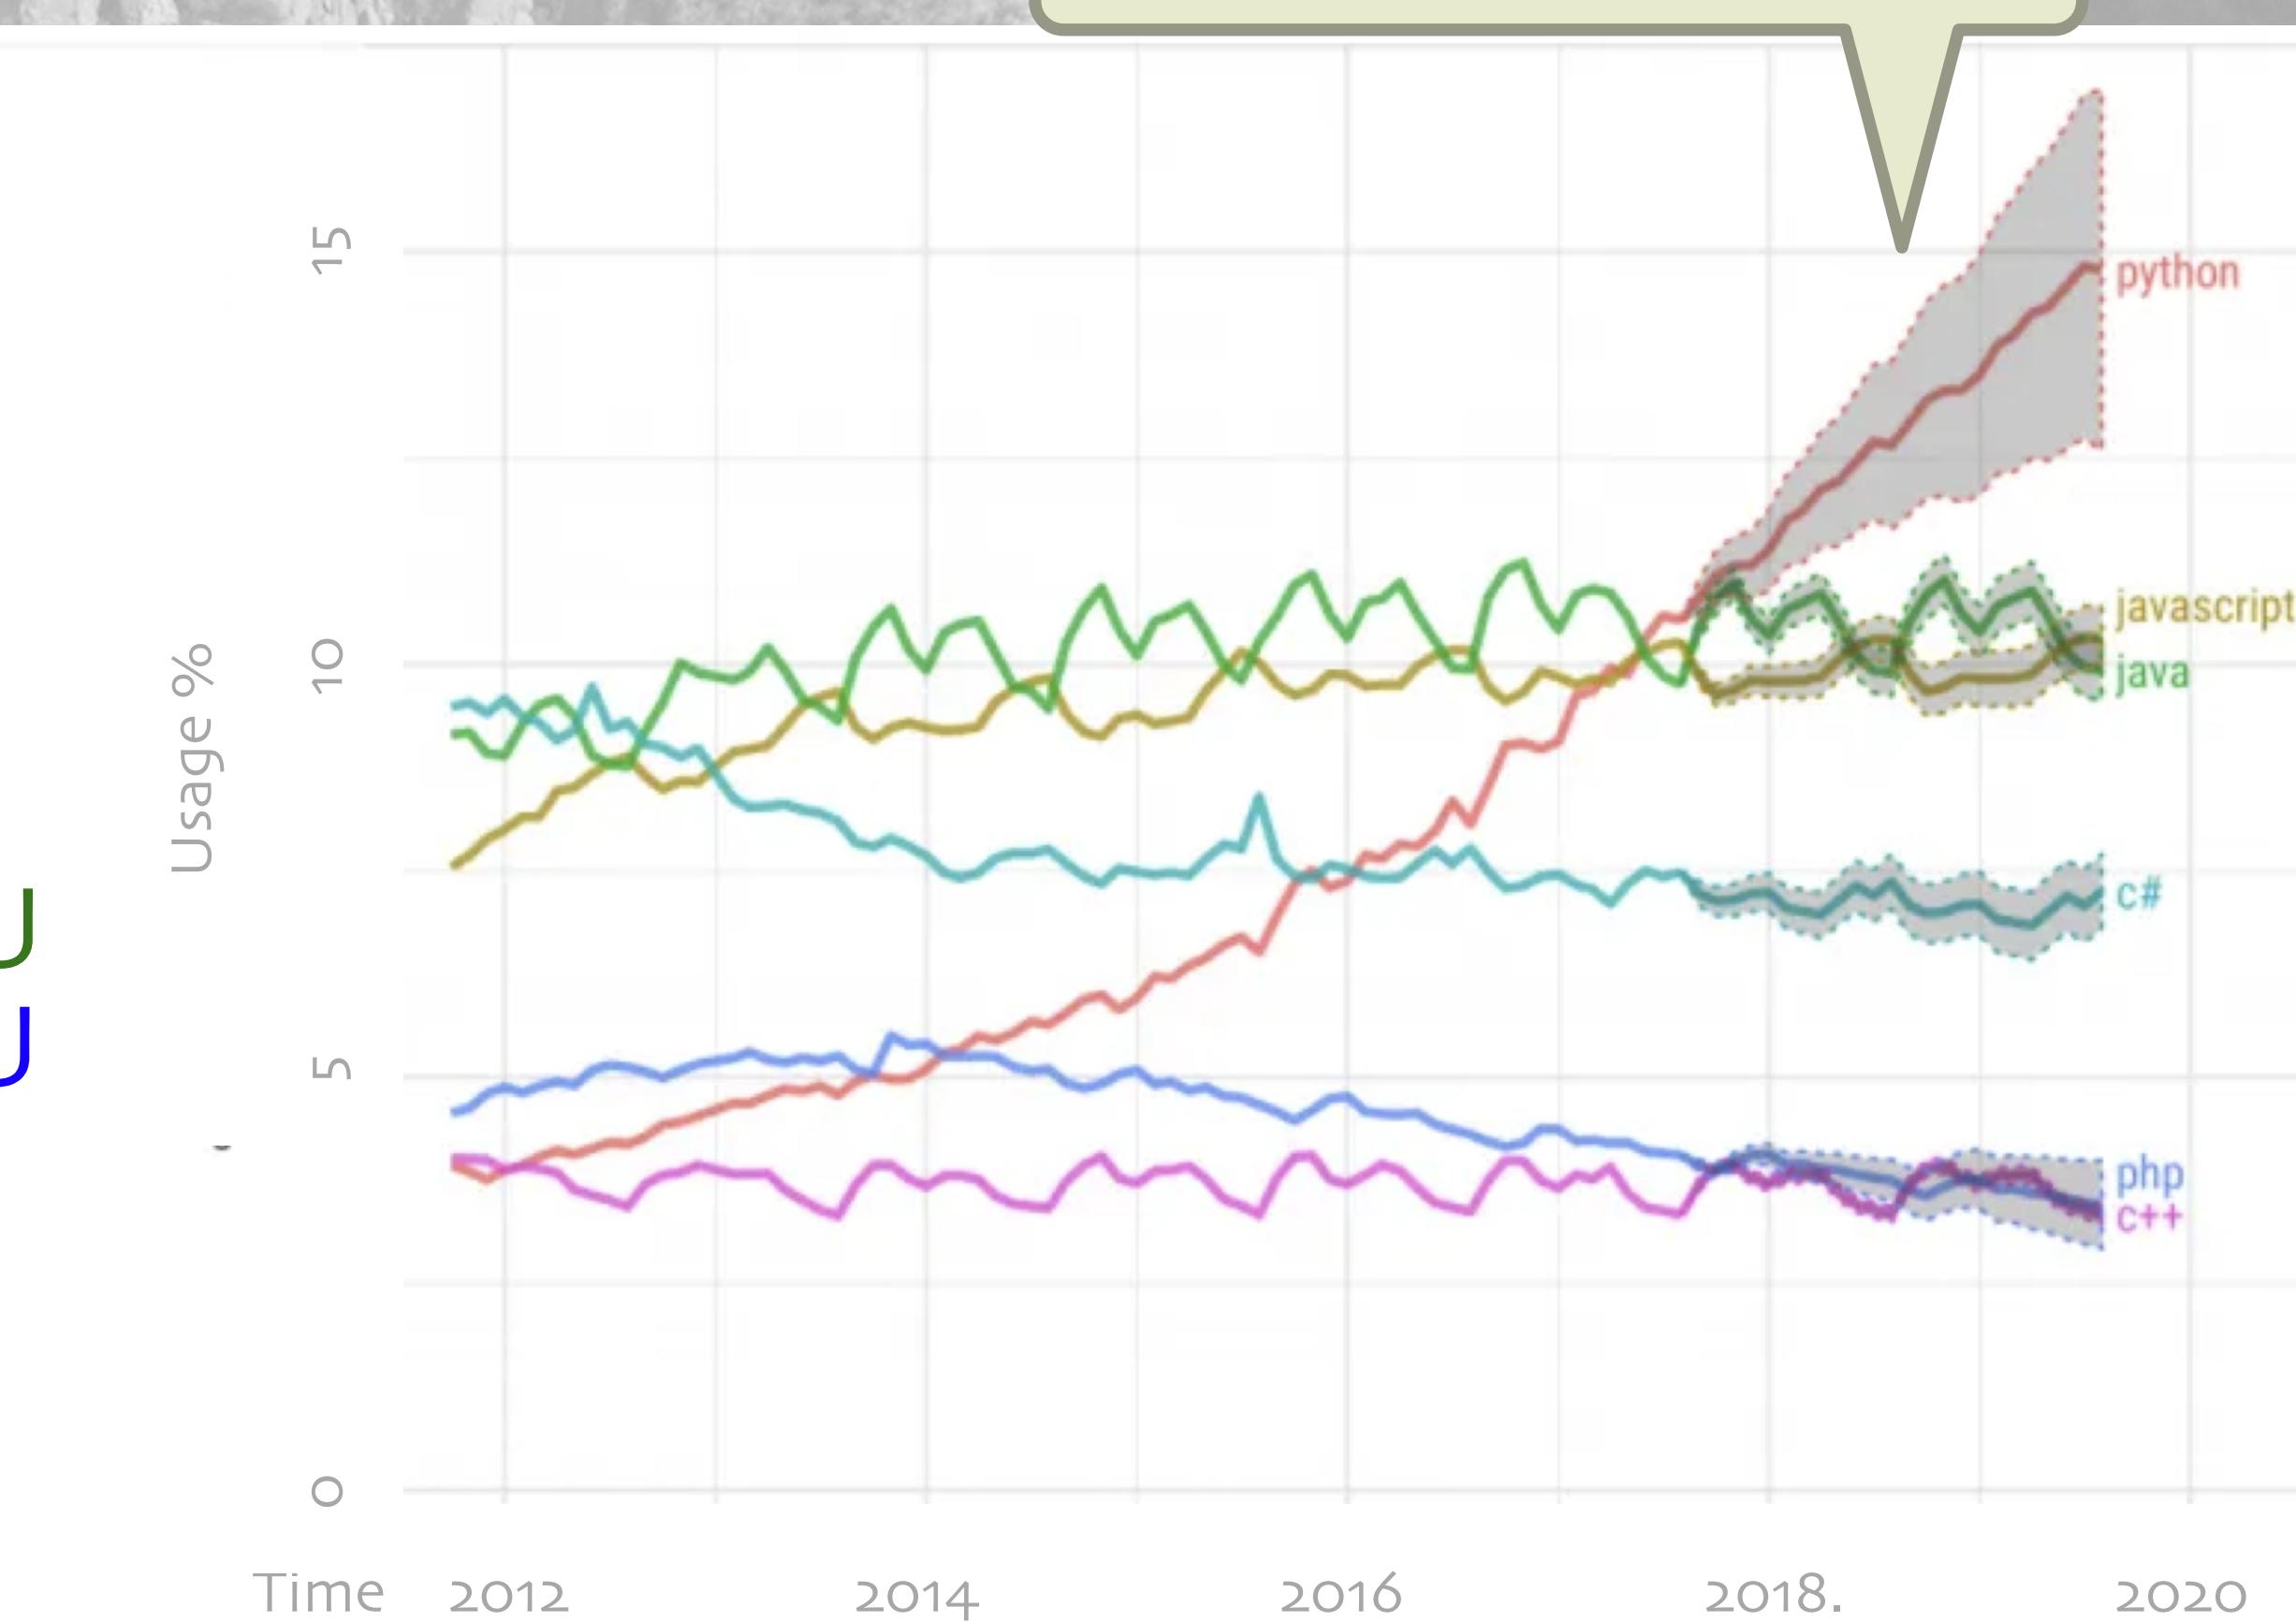
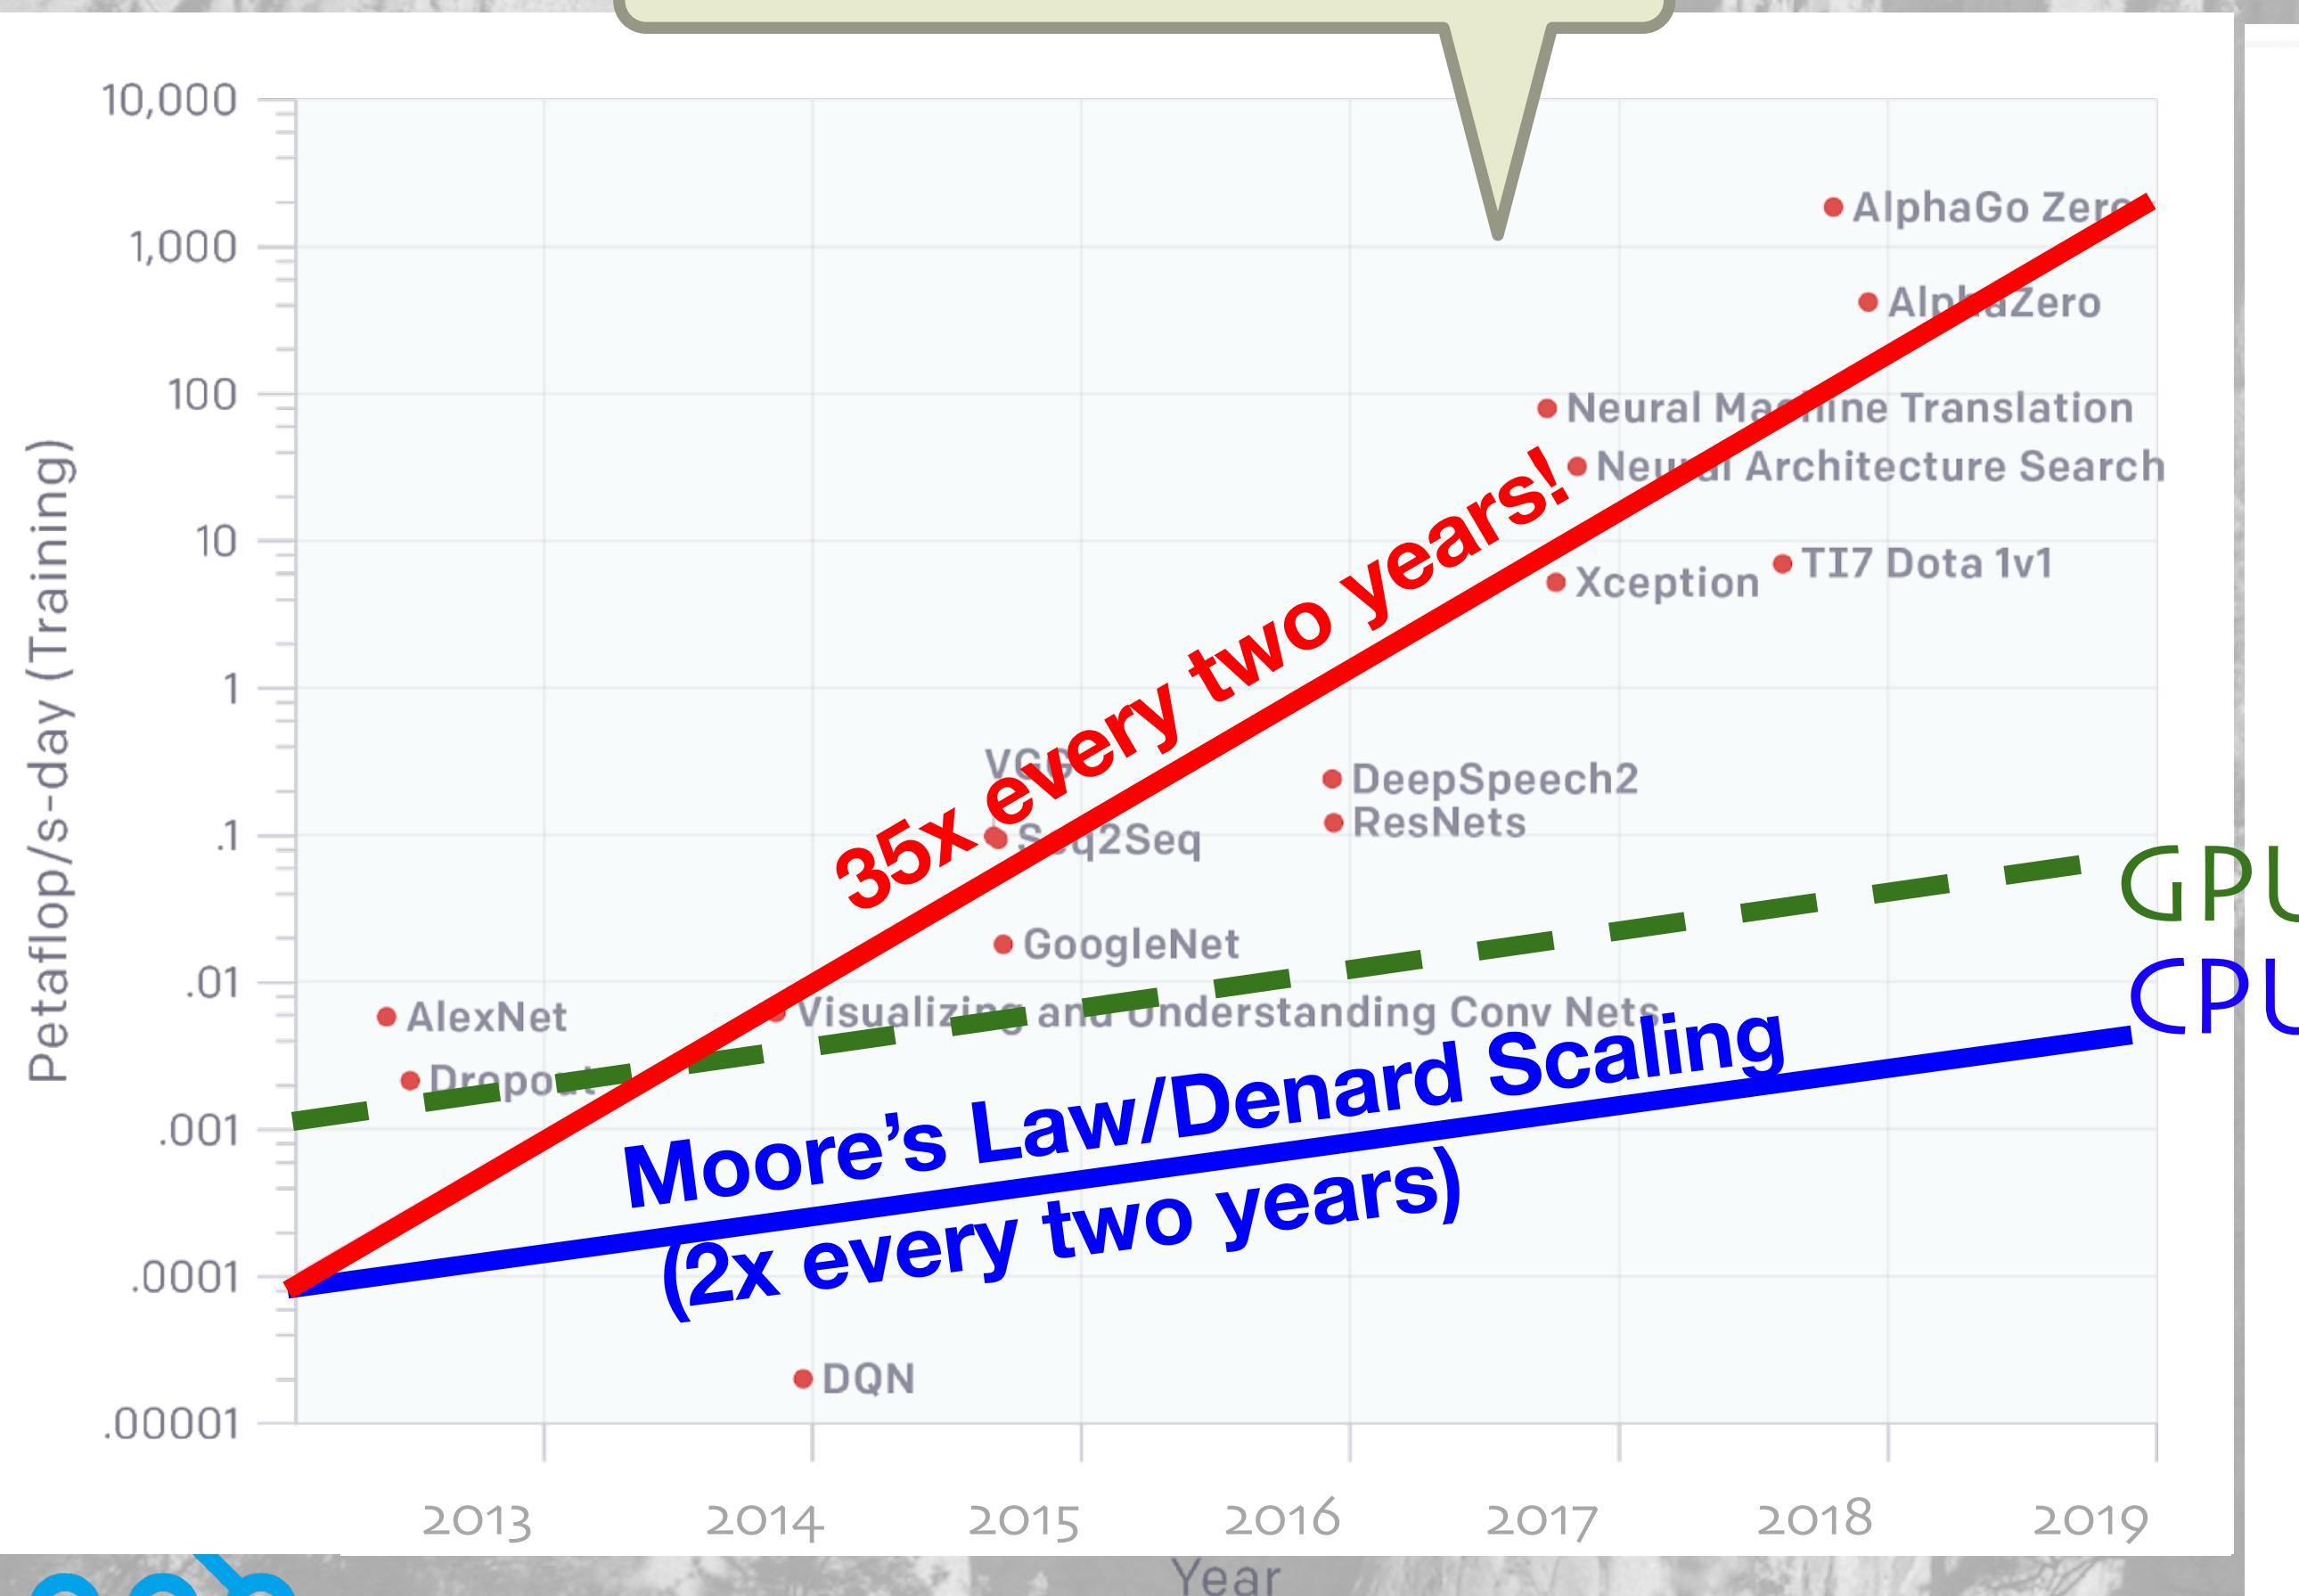


Two Major Trends

Model sizes and therefore compute requirements outstripping Moore's Law

Hence, there is a pressing need for robust, easy to use solutions for distributed Python

Python growth driven by ML/AI and other data science workloads



The ML Landscape Today

All require distributed implementations to scale

Featurization



Streaming



Hyperparam
Tuning



Training



Simulation



Model
Serving



The Ray Vision: Sharing a Common Framework

Featurization

Streaming

Hyperparam
Tuning

Training

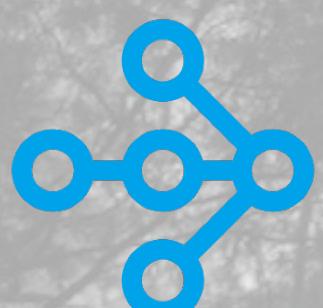
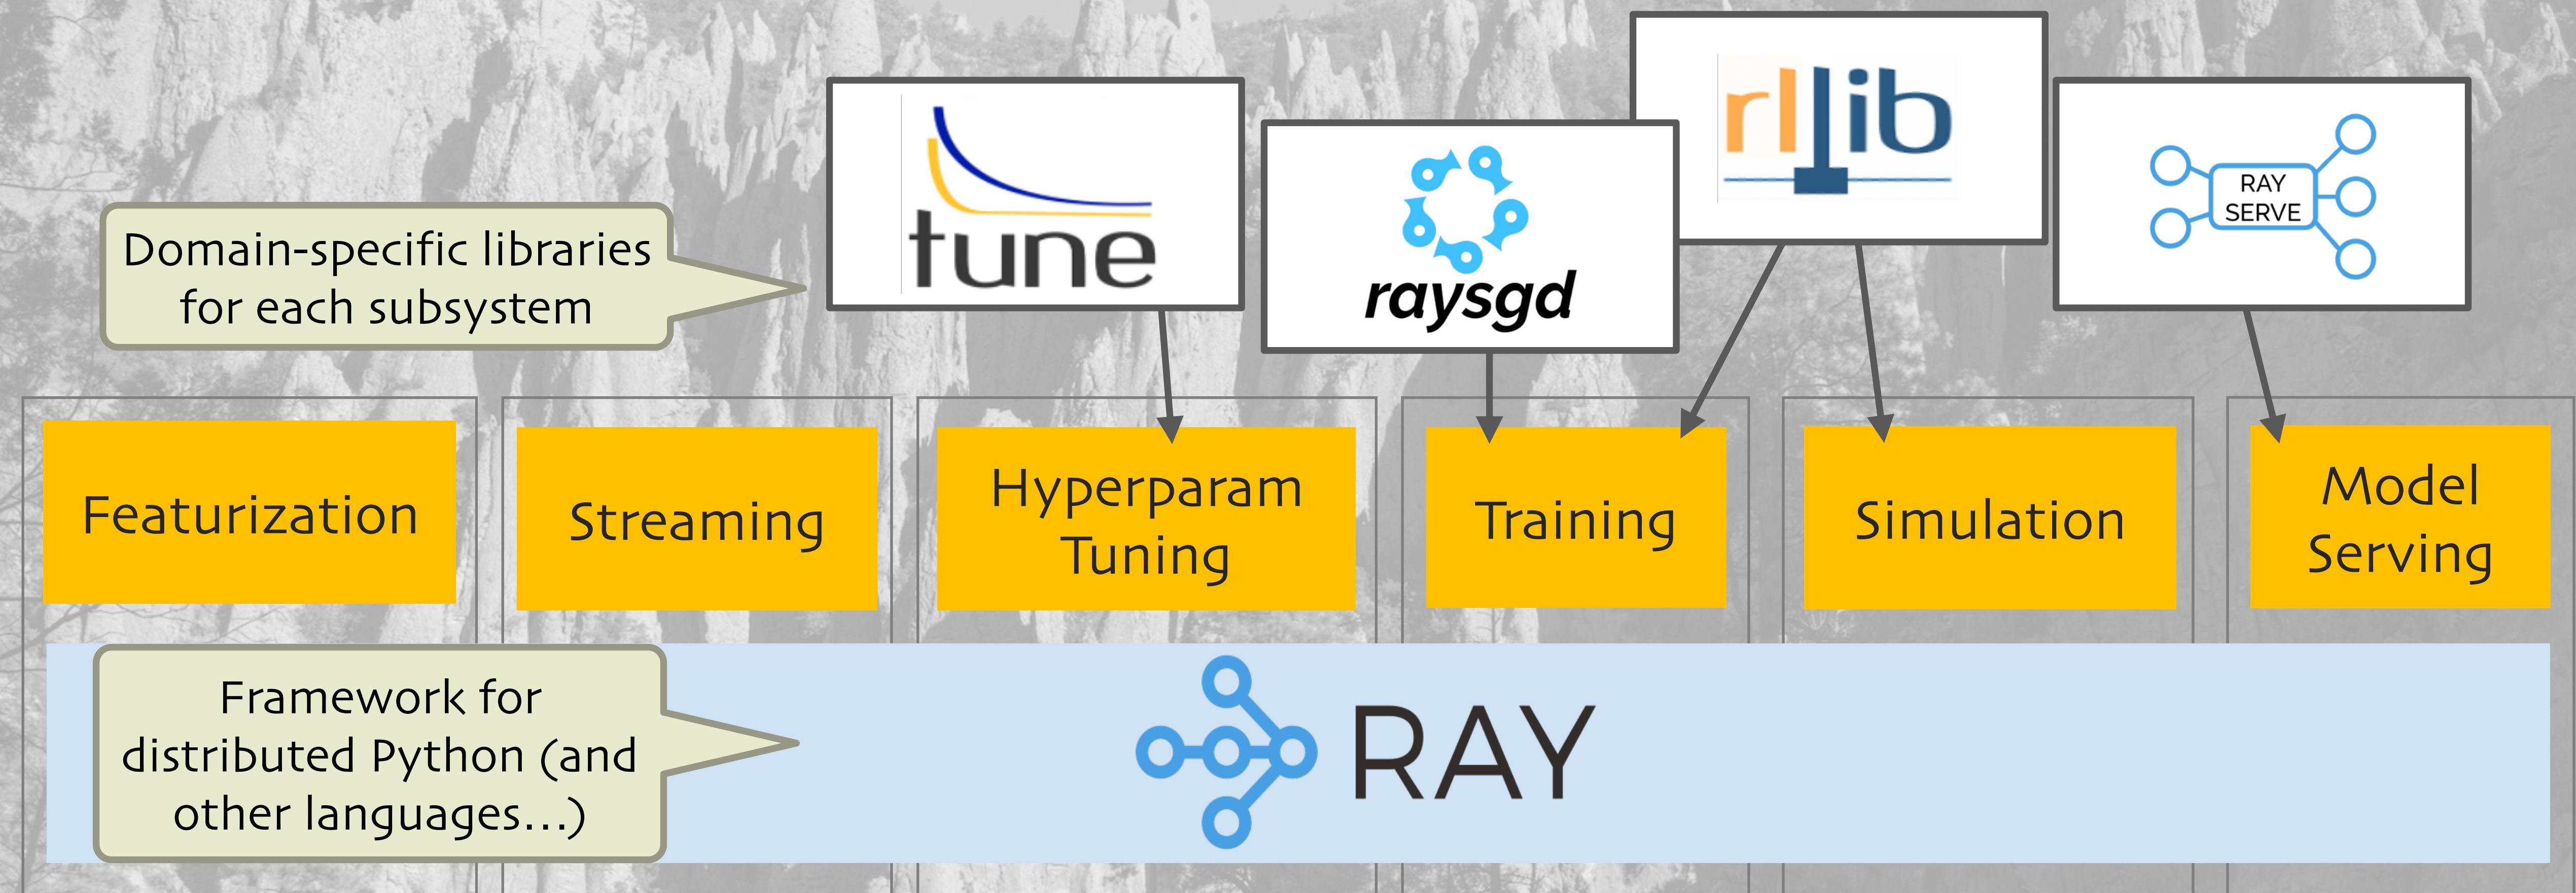
Simulation

Model
Serving

Framework for
distributed Python (and
other languages...)



The Ray Vision: Sharing a Common Framework



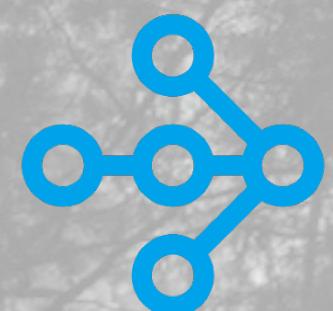
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
def add_arrays(a, b):  
    return np.add(a, b)
```

The Python you
already know...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

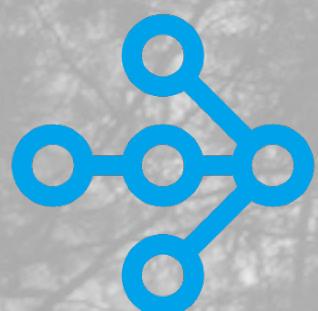
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

For completeness, add these first:

```
import ray  
import numpy as np  
ray.init()
```

Now these functions
are remote "tasks"



API - Designed to Be Intuitive and Concise

Functions -> Tasks

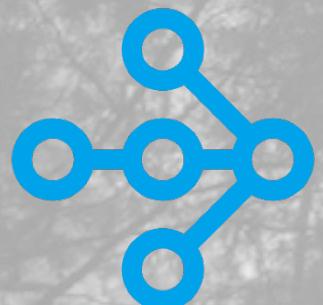
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
ref1 = make_array.remote(...)
```

make_array

ref1



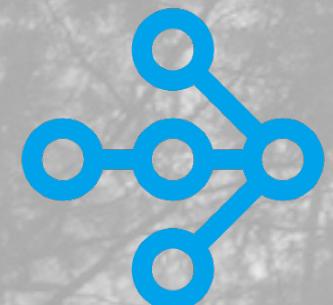
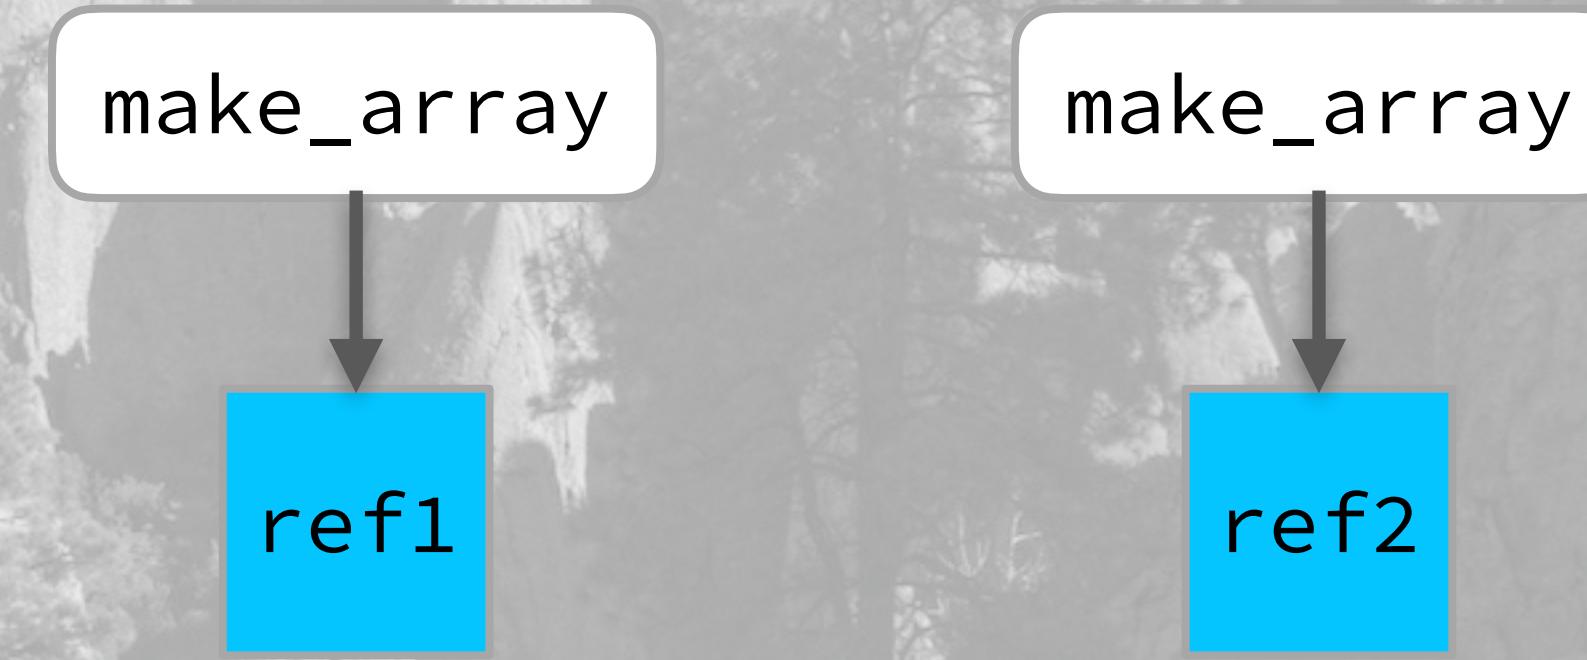
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)
```



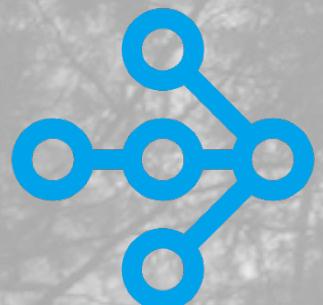
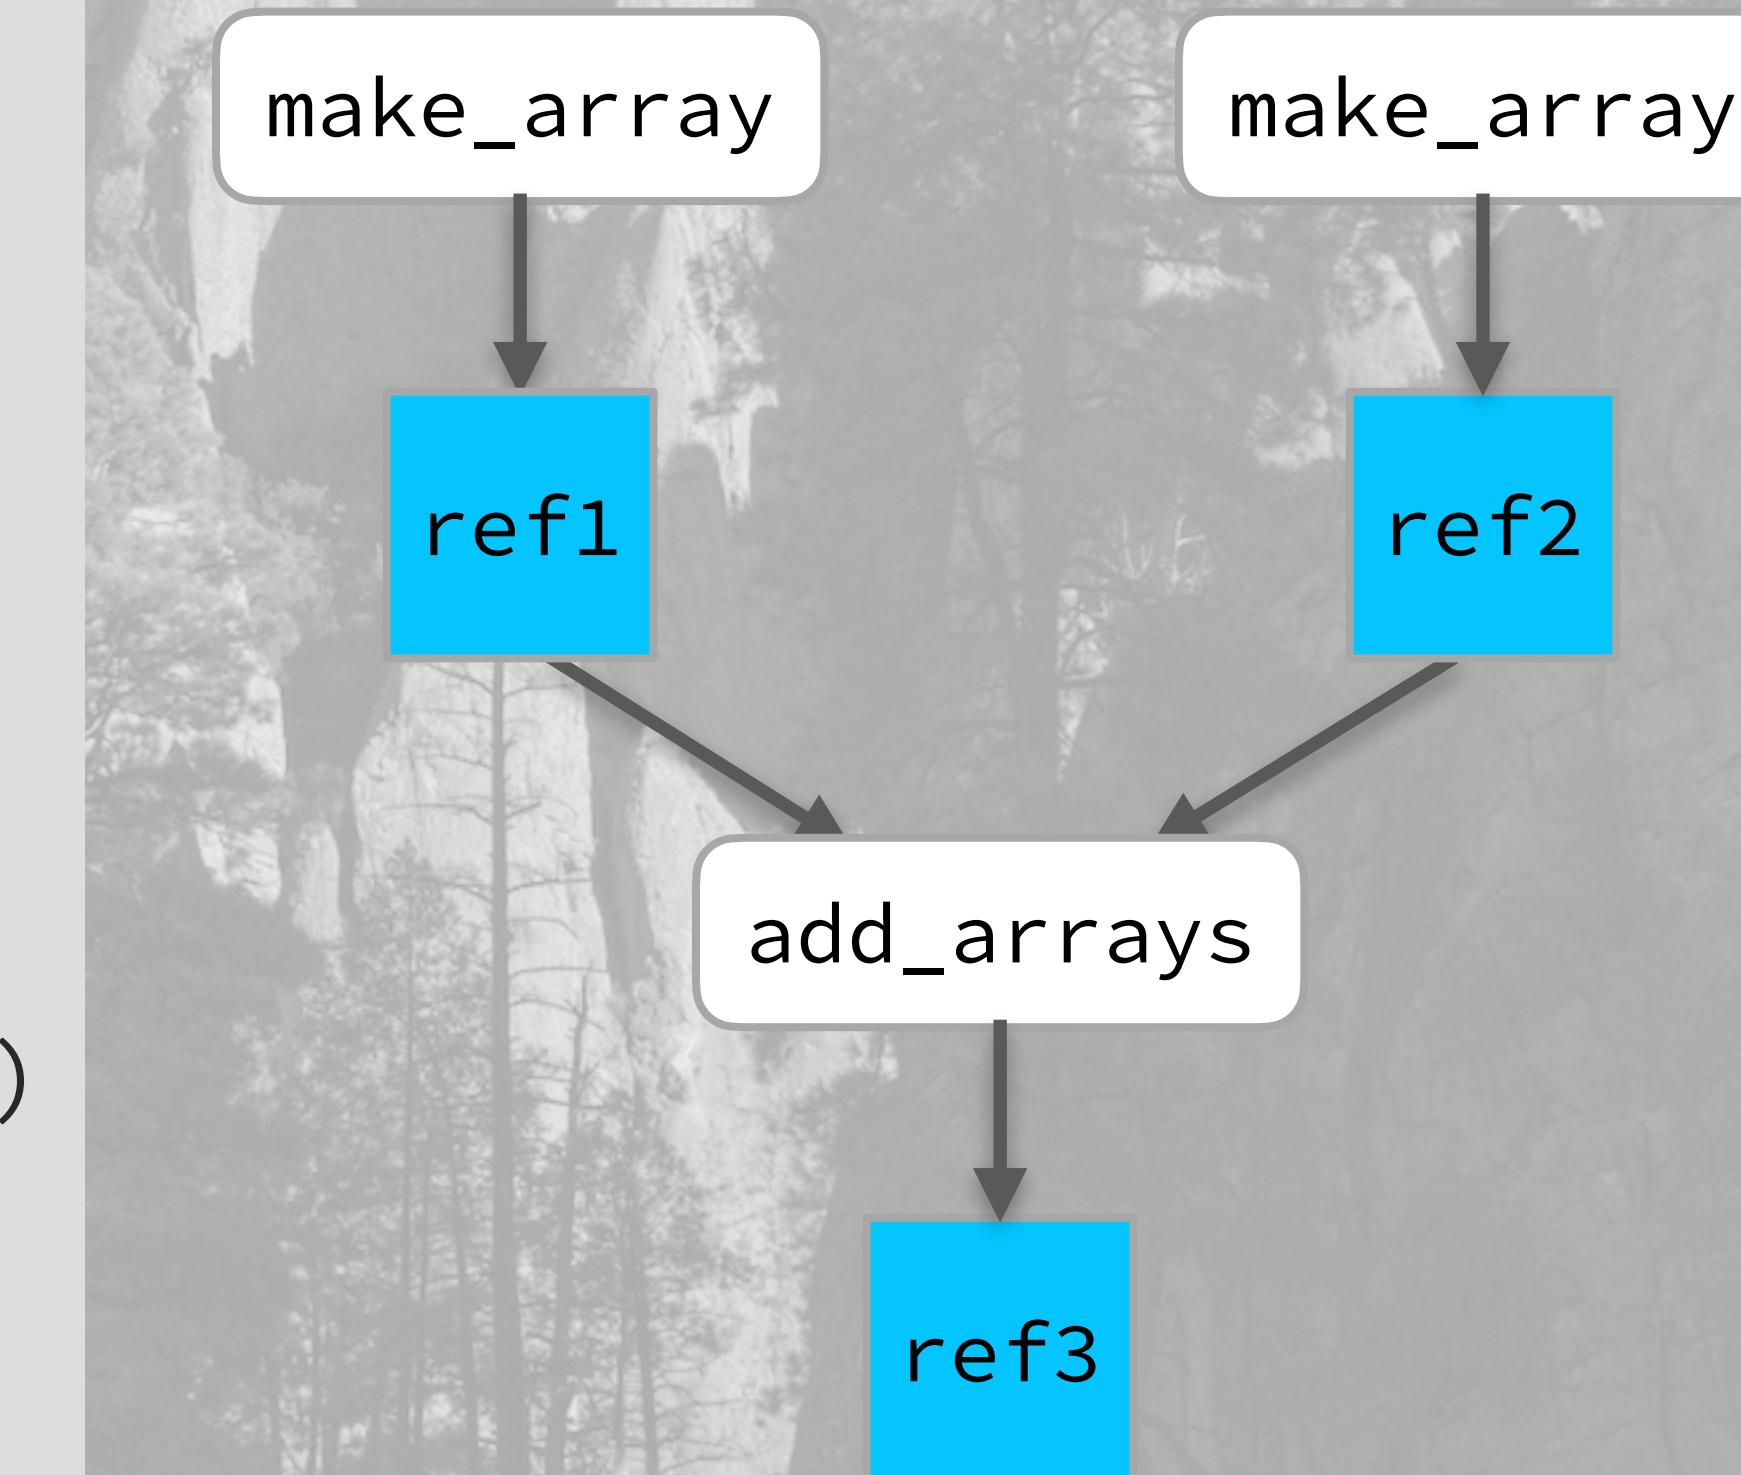
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

ref1 = make_array.remote(...)
ref2 = make_array.remote(...)
ref3 = add_arrays.remote(ref1, ref2)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

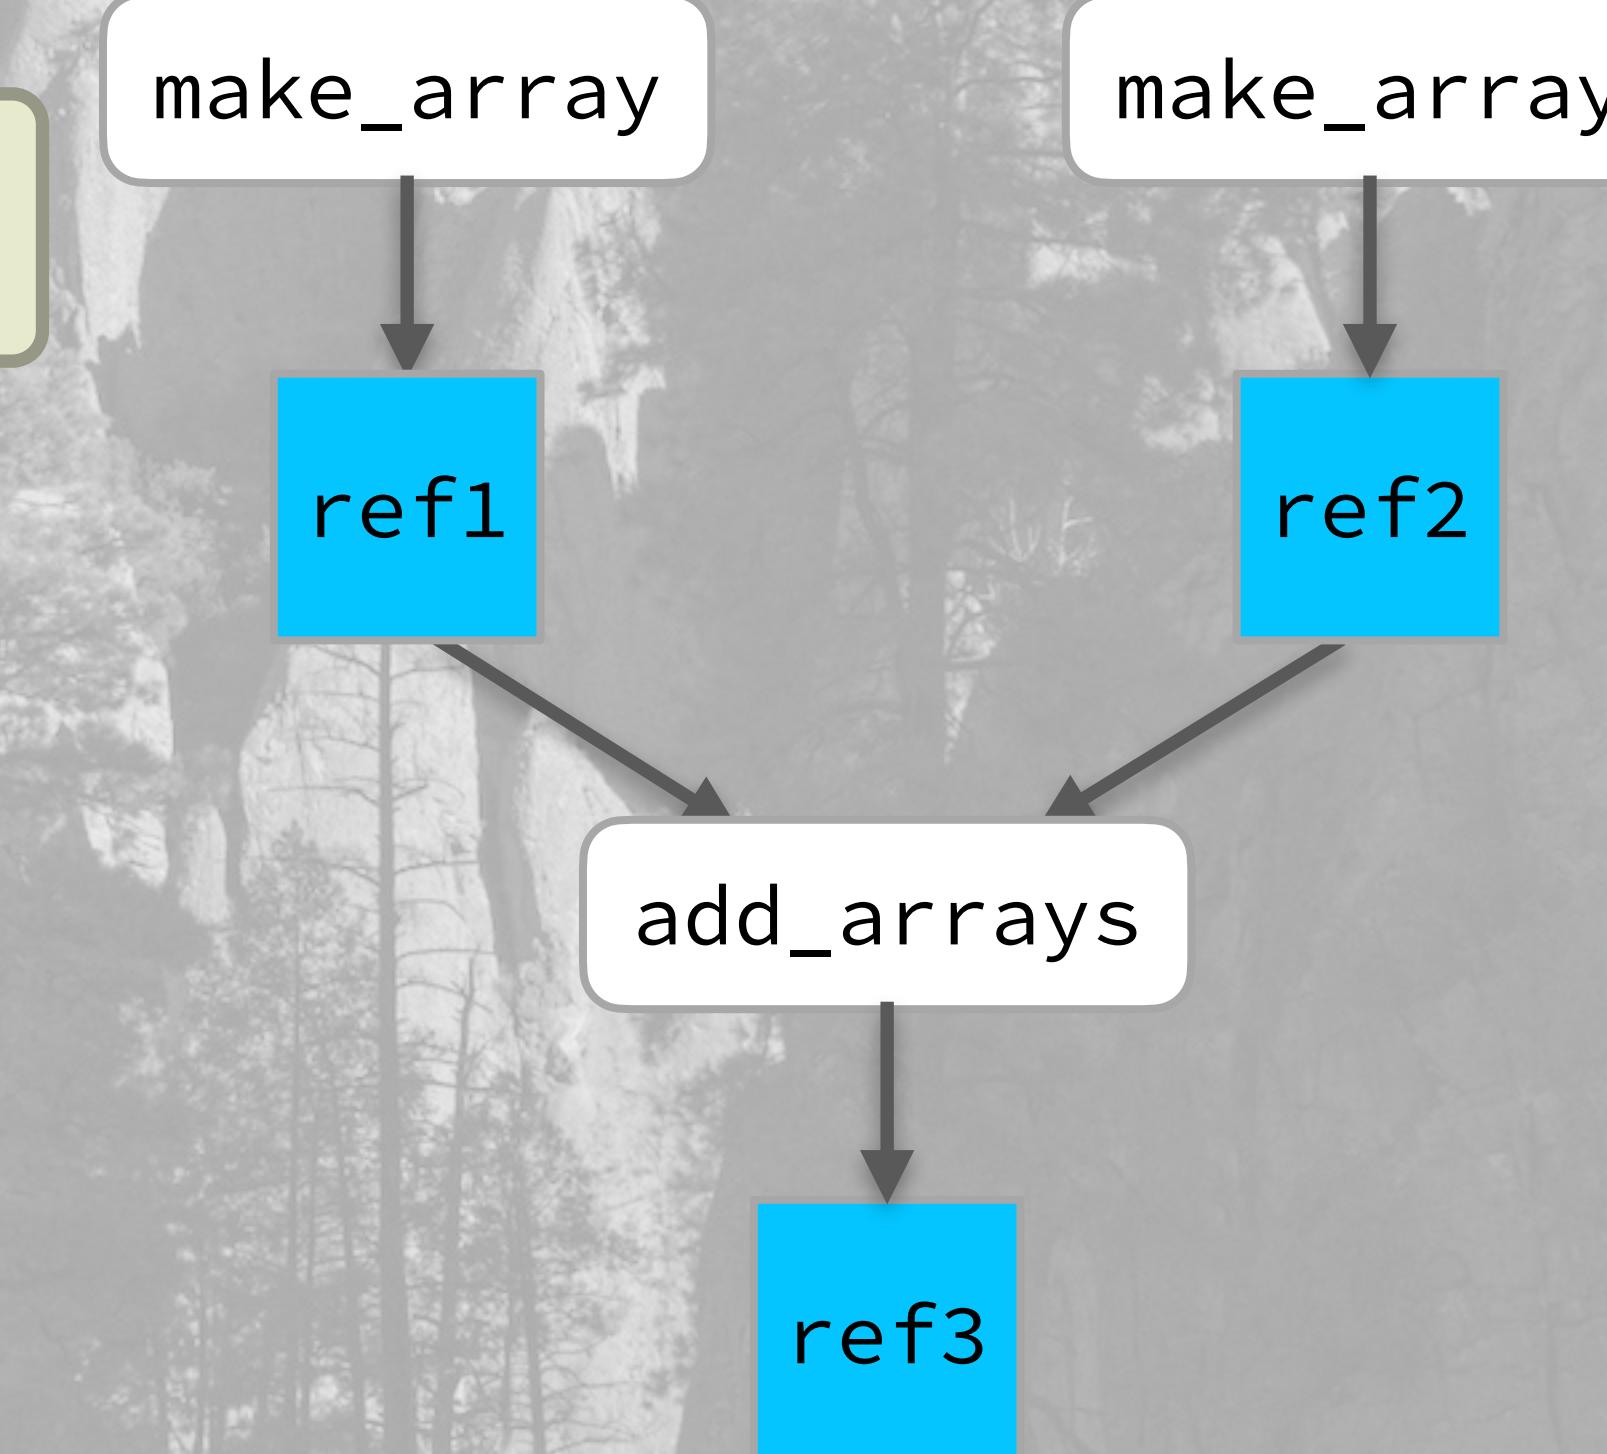
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a
```

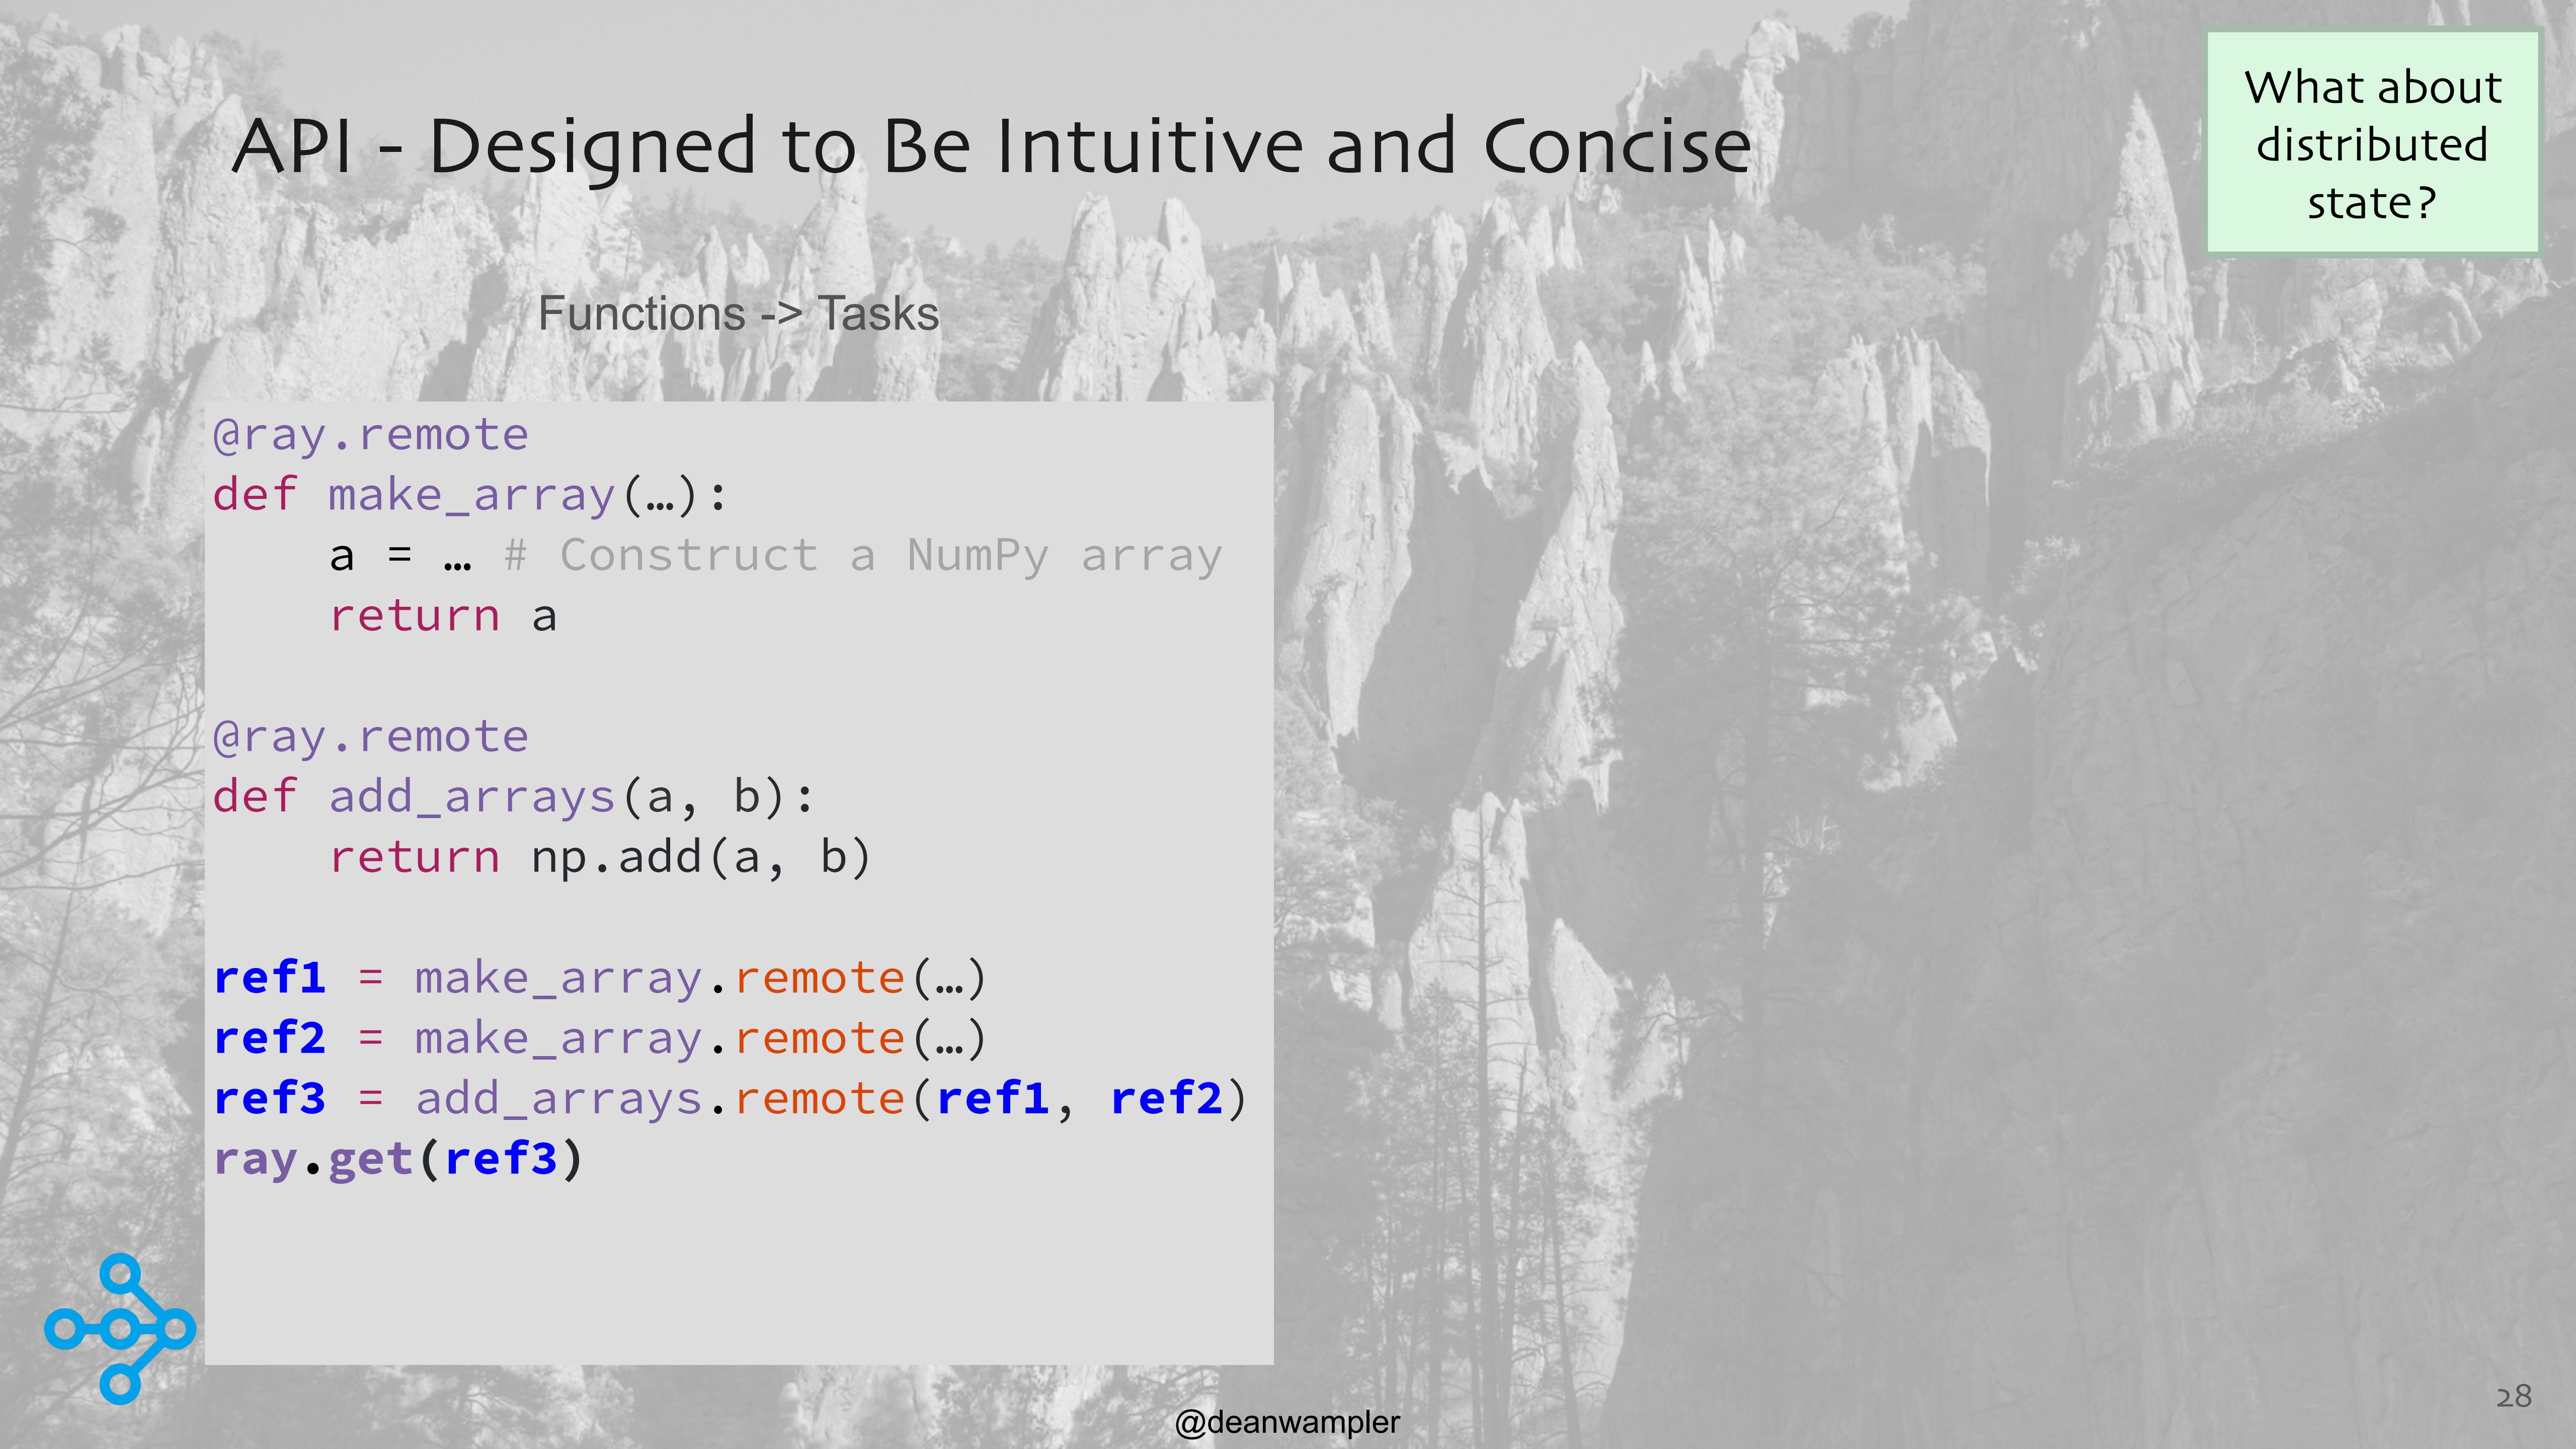
```
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)
```

```
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```

Ray handles sequencing
of async dependencies

Ray handles extracting the
arrays from the object refs





What about
distributed
state?

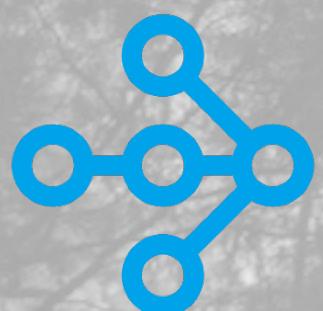
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

ref1 = make_array.remote(...)
ref2 = make_array.remote(...)
ref3 = add_arrays.remote(ref1, ref2)
ray.get(ref3)
```



API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

ref1 = make_array.remote(...)
ref2 = make_array.remote(...)
ref3 = add_arrays.remote(ref1, ref2)
ray.get(ref3)
```

Classes -> Actors

```
class Counter(object):
    def __init__(self):
        self.value = 0
    def increment(self):
        self.value += 1
    return self.value
```

The Python
classes you
love...



API - Designed to Be Intuitive and Concise

Functions -> Tasks

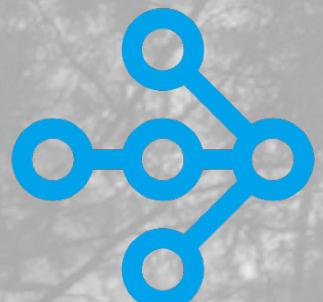
```
@ray.remote  
def make_array(...):  
    a = ... # Construct a NumPy array  
    return a  
  
@ray.remote  
def add_arrays(a, b):  
    return np.add(a, b)  
  
ref1 = make_array.remote(...)  
ref2 = make_array.remote(...)  
ref3 = add_arrays.remote(ref1, ref2)  
ray.get(ref3)
```

Classes -> Actors

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def increment(self):  
        self.value += 1  
        return self.value  
    def get_count(self):  
        return self.value
```

... now a remote
“actor”

You need a
“getter” method
to read the state.



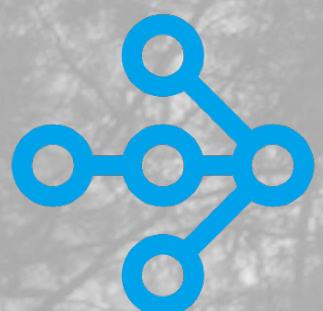
API - Designed to Be Intuitive and Concise

Functions -> Tasks

```
@ray.remote
def make_array(...):
    a = ... # Construct a NumPy array
    return a

@ray.remote
def add_arrays(a, b):
    return np.add(a, b)

ref1 = make_array.remote(...)
ref2 = make_array.remote(...)
ref3 = add_arrays.remote(ref1, ref2)
ray.get(ref3)
```



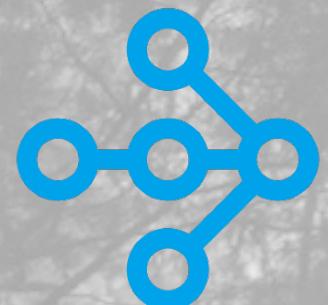
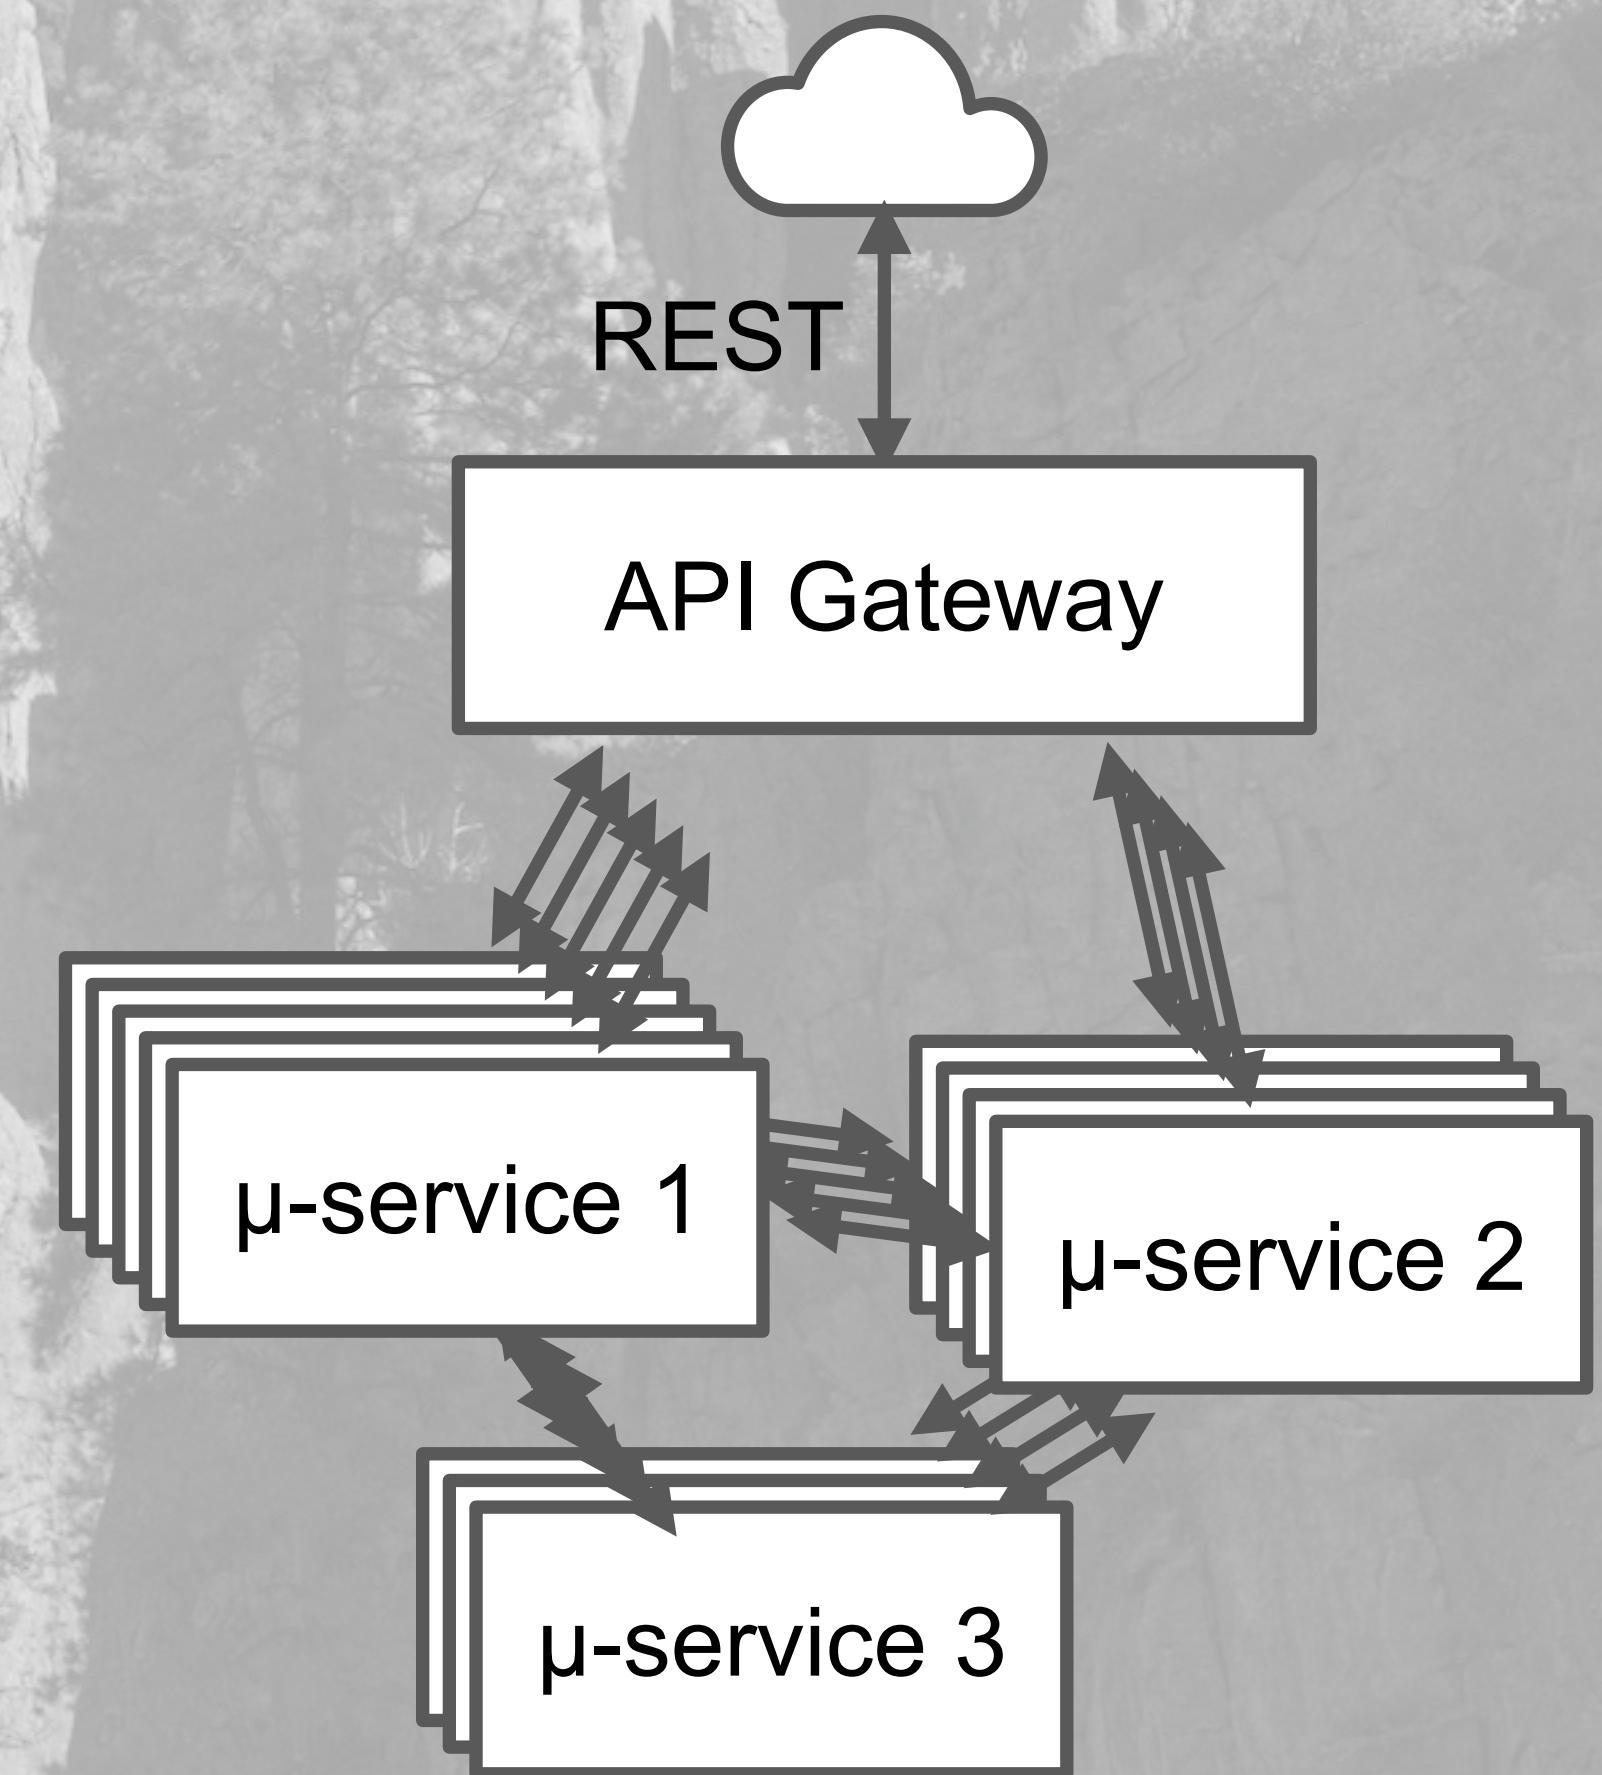
Classes -> Actors

```
@ray.remote
class Counter(object):
    def __init__(self):
        self.value = 0
    def increment(self):
        self.value += 1
        return self.value
    def get_count(self):
        return self.value

c = Counter.remote()
ref4 = c.increment.remote()
ref5 = c.increment.remote()
ray.get([ref4, ref5]) # [1, 2]
```

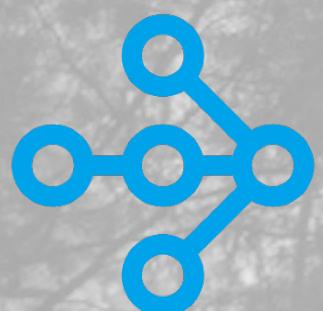
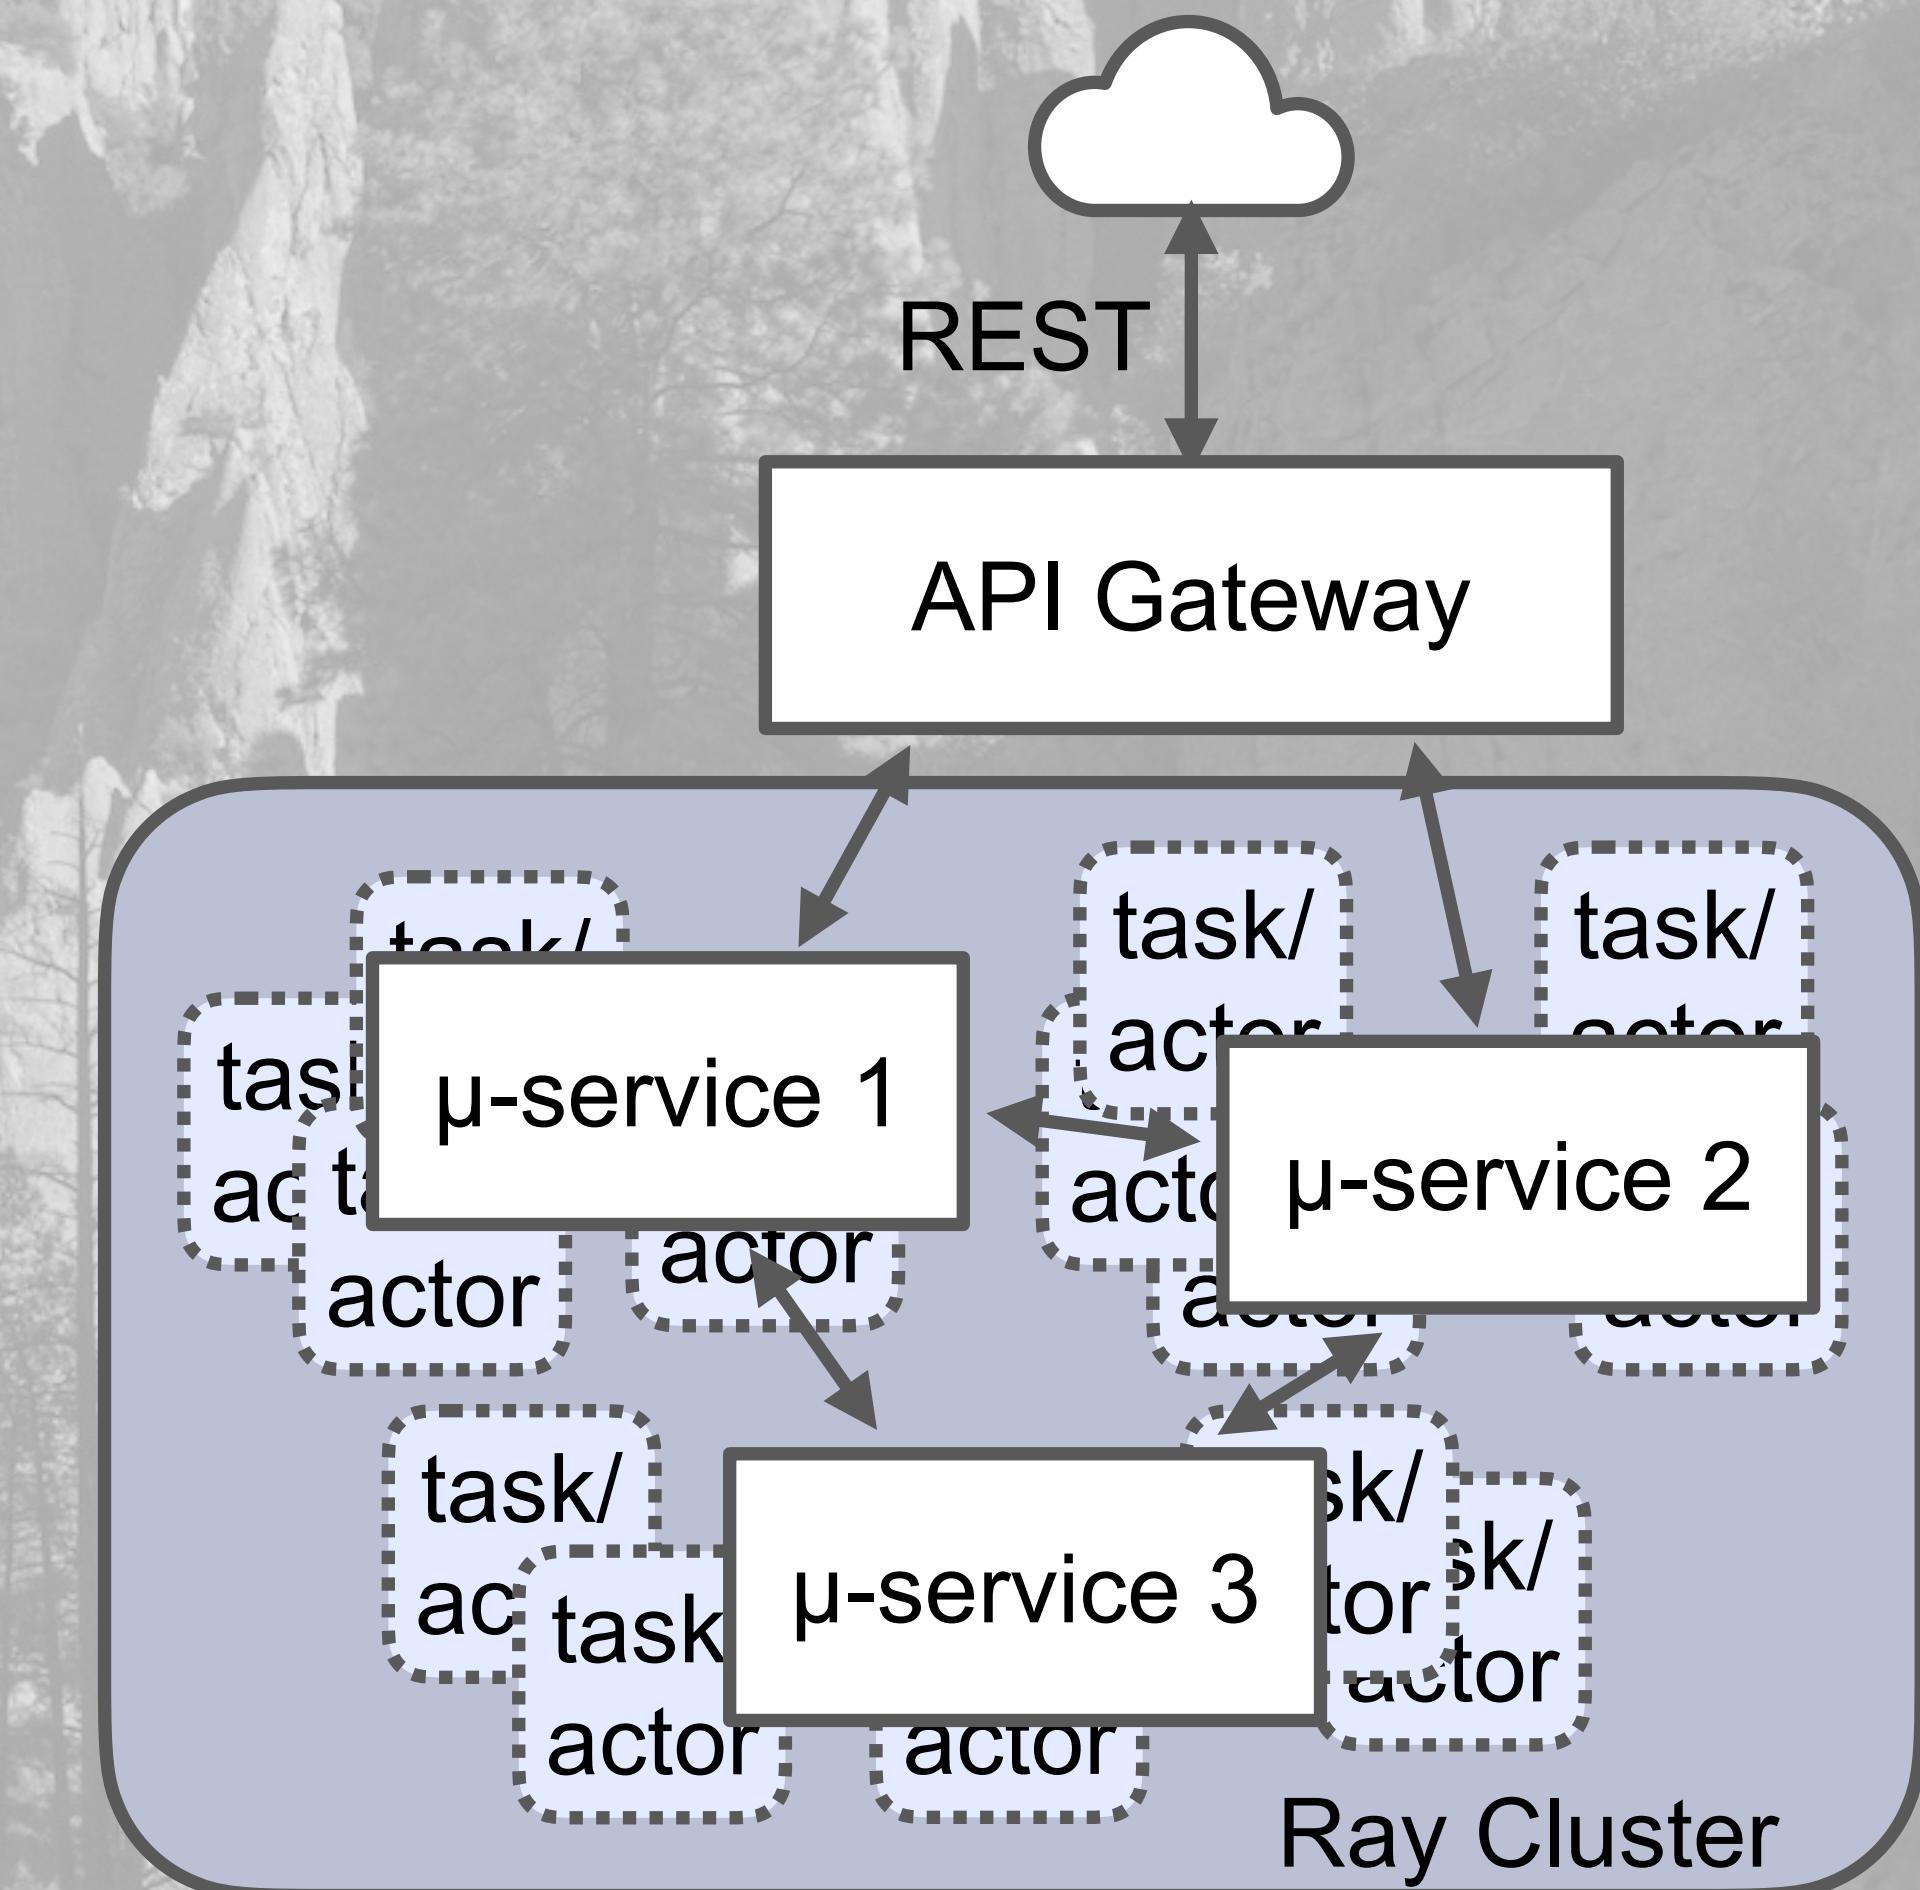
Another View - Distributed ML/AI as Microservices?

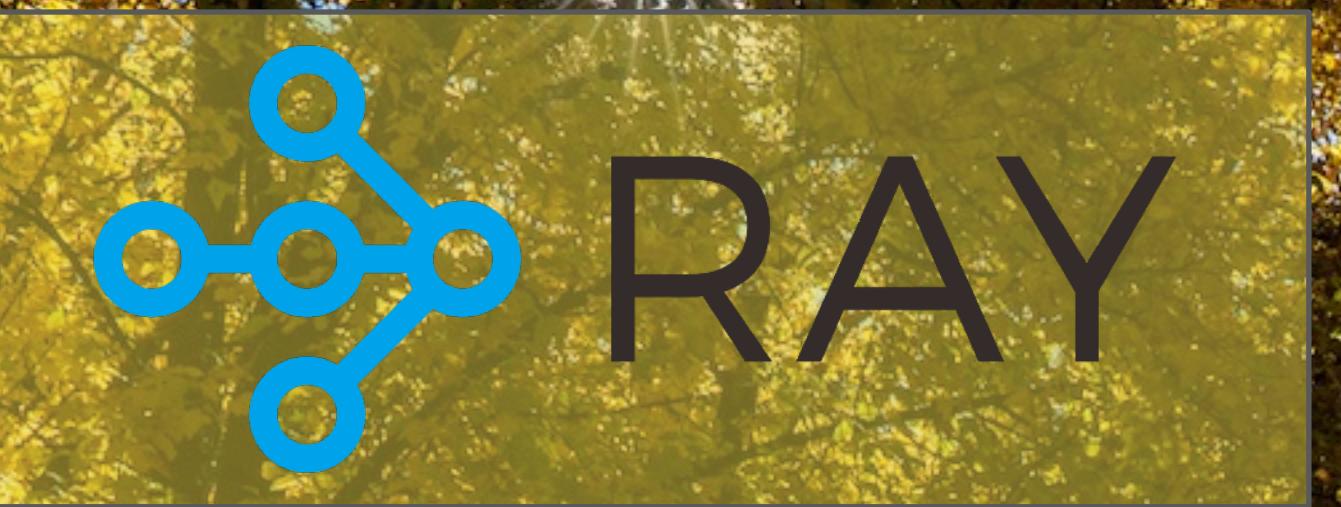
- Each microservice has a different number of instances for scalability and resiliency
- But they have to be managed **explicitly**



Microservices - Simplified

- With Ray, you have one “logical” instance to manage and Ray does the cluster-wide scaling for you.



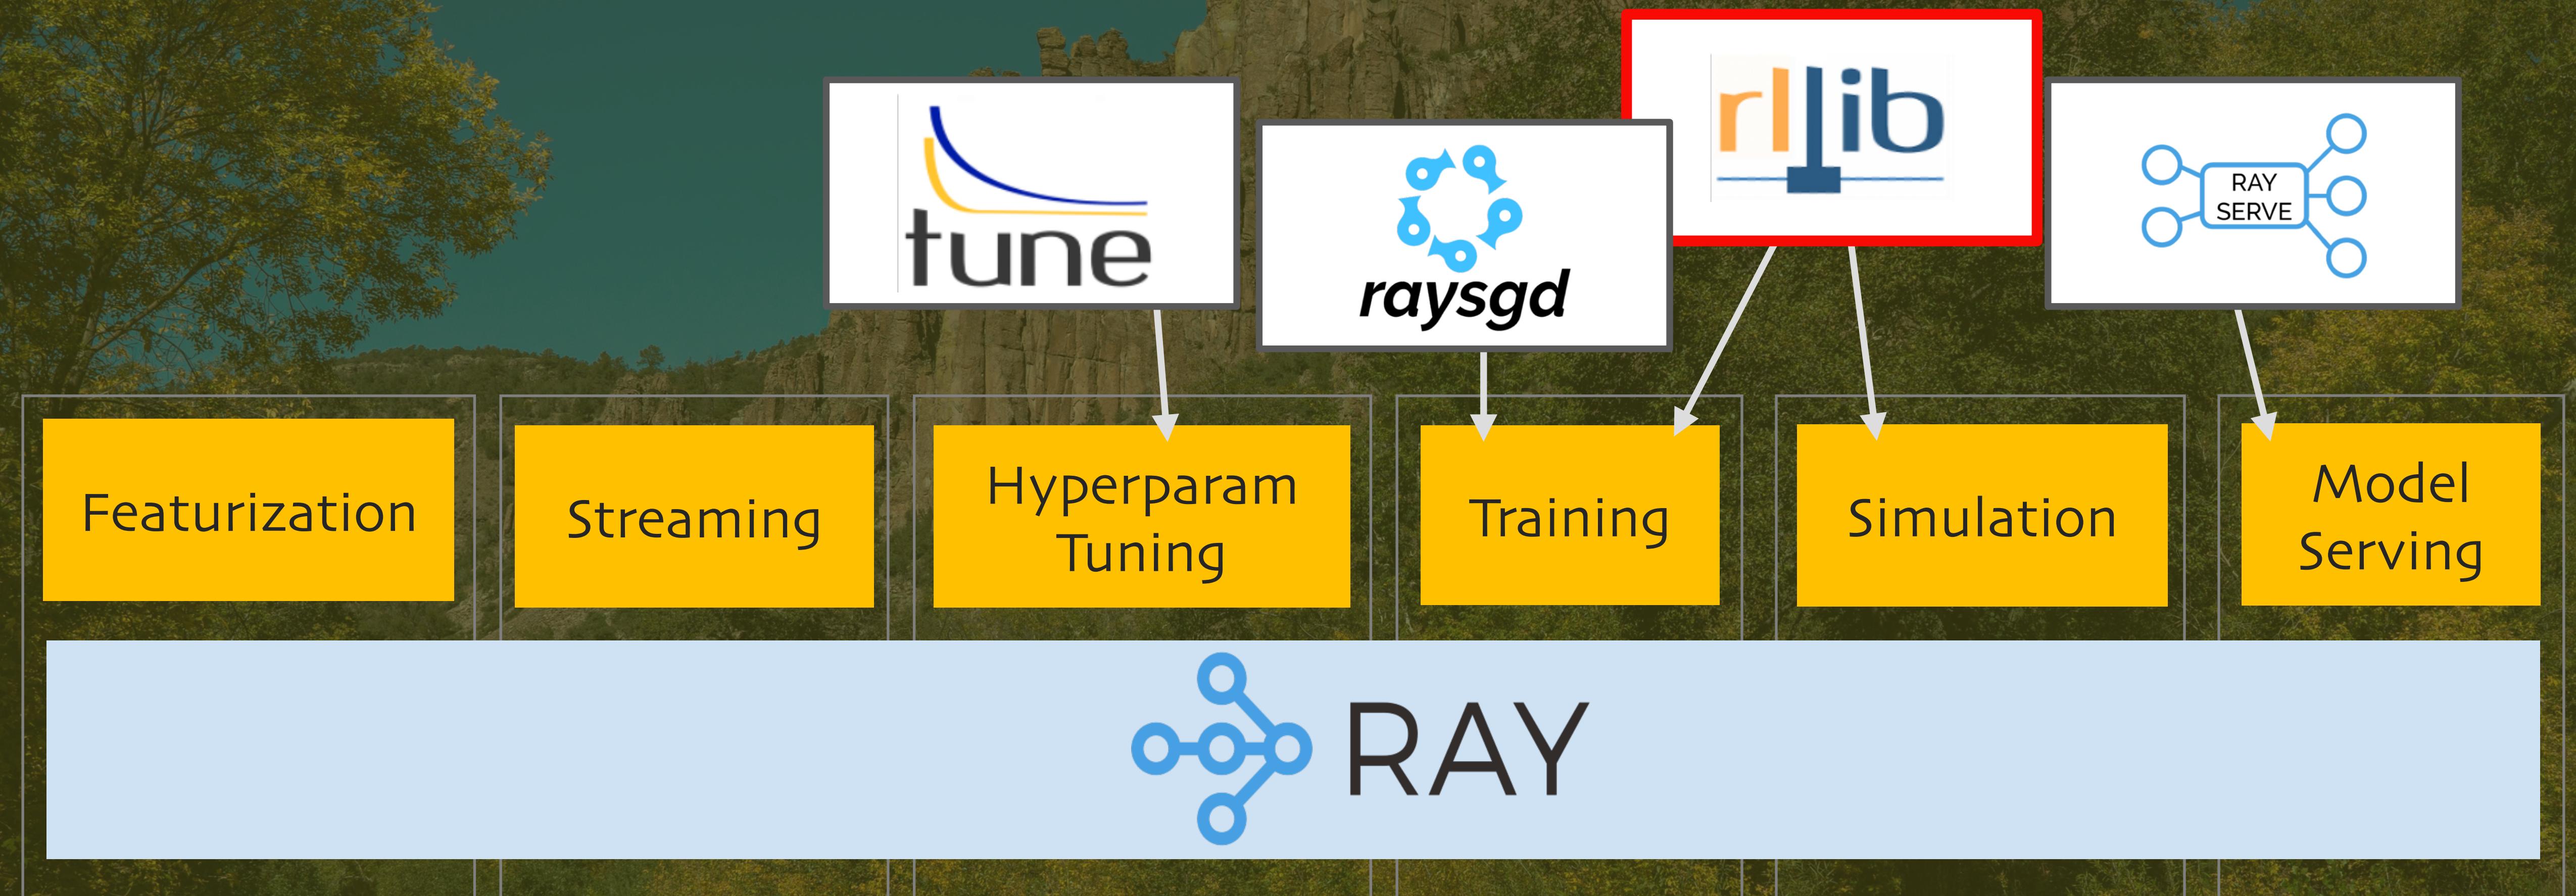


Ray RLlib



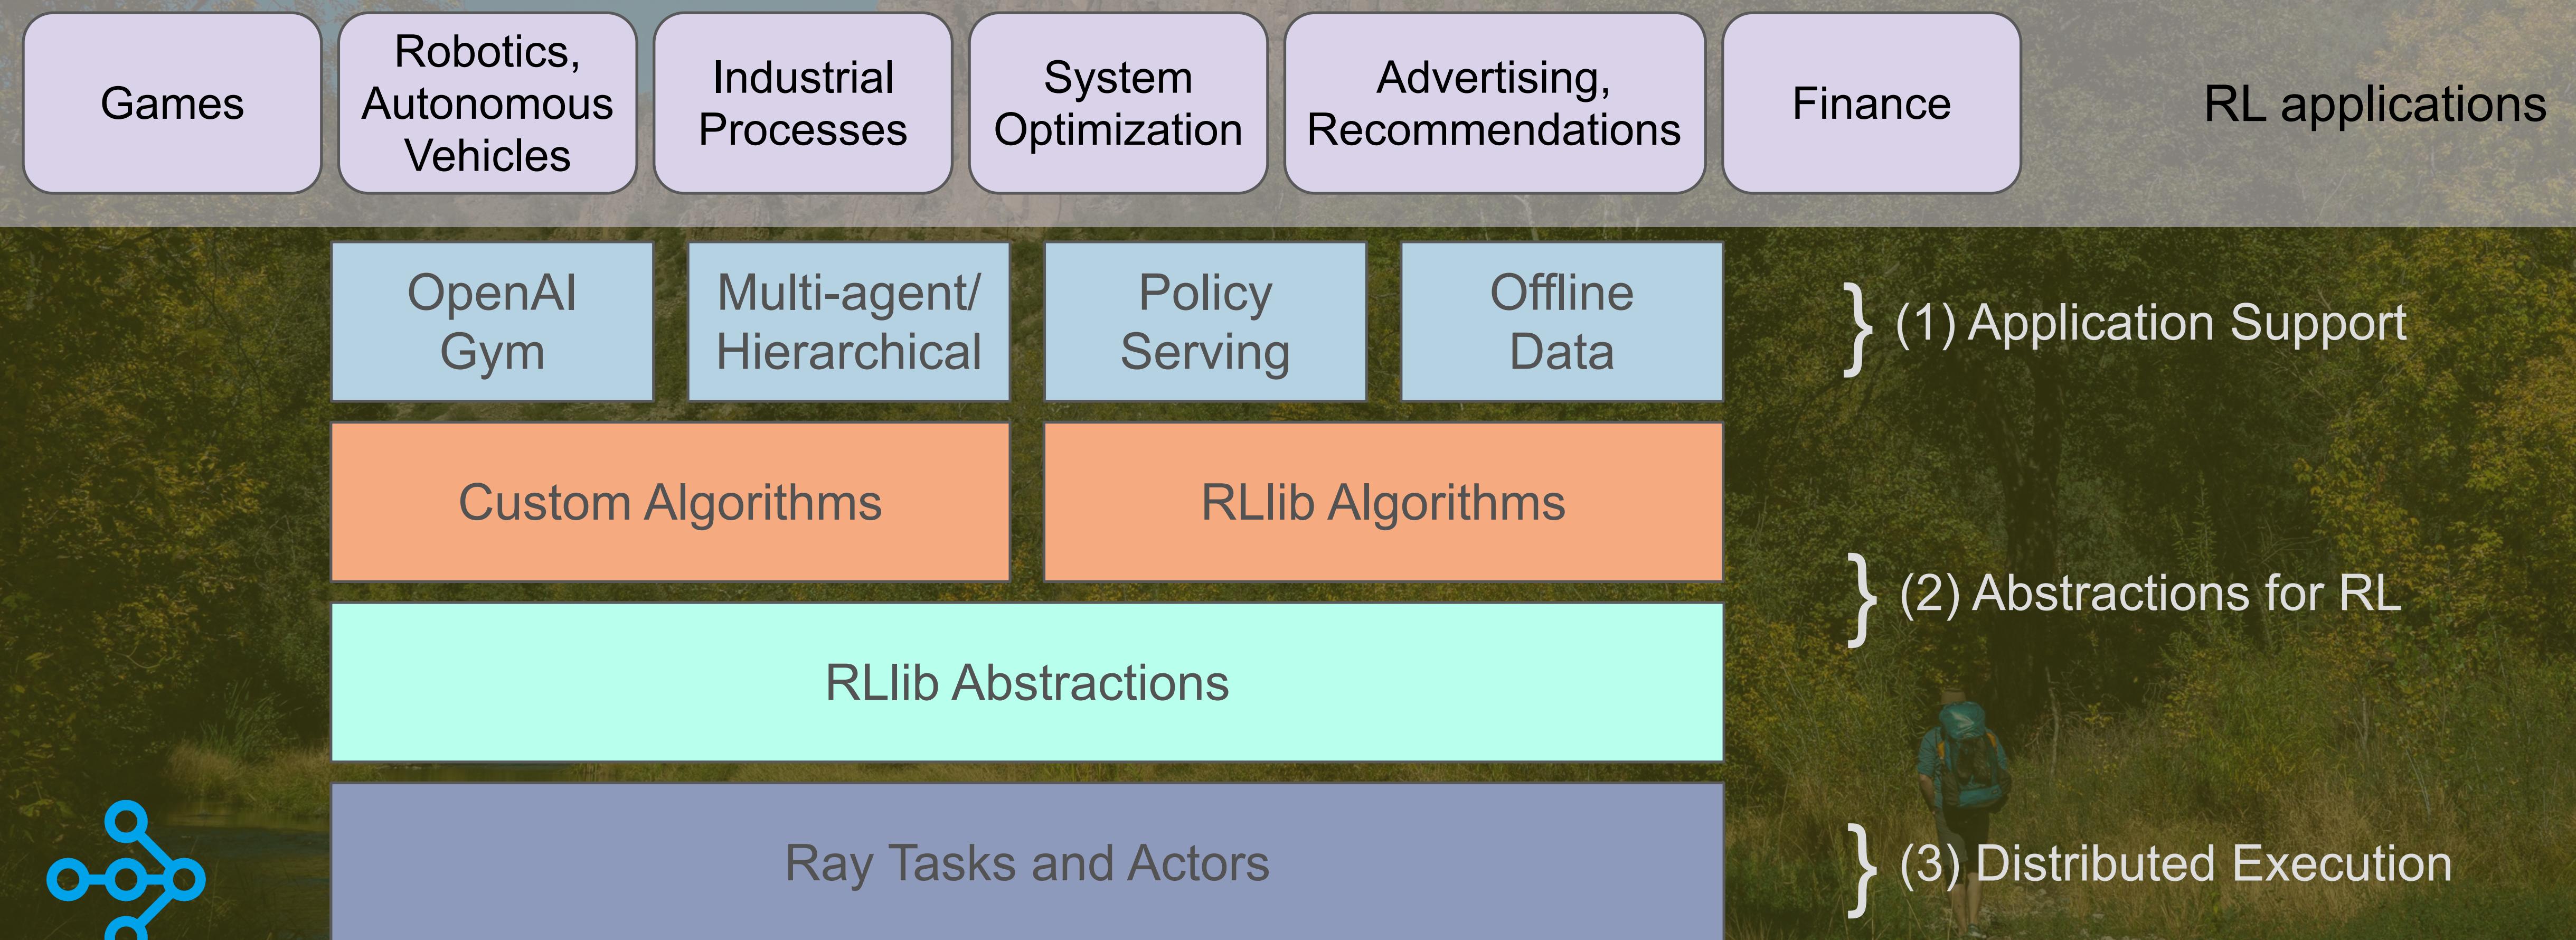
@deanwampler

Reinforcement Learning - Ray RLlib



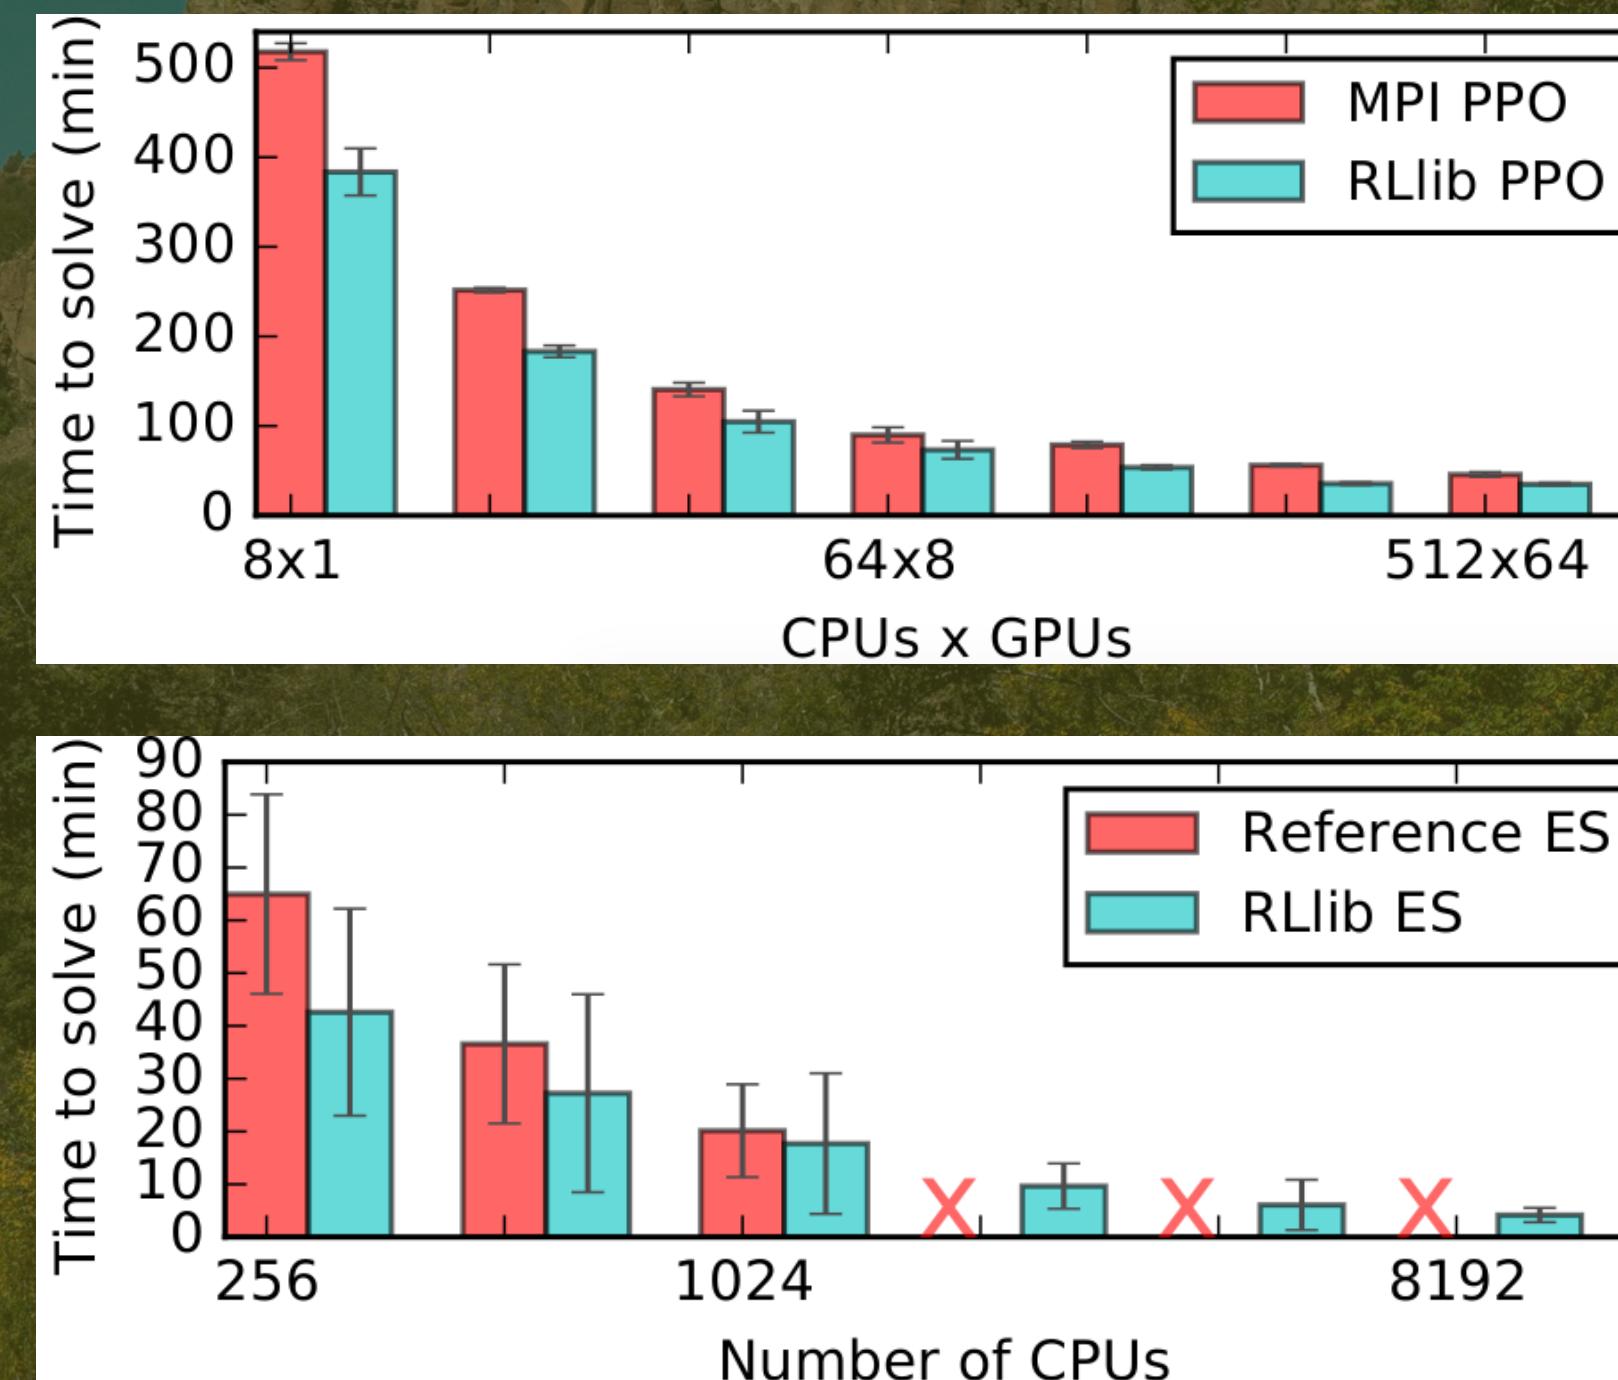
rllib.io

RLLib: A Scalable, Unified Library for RL



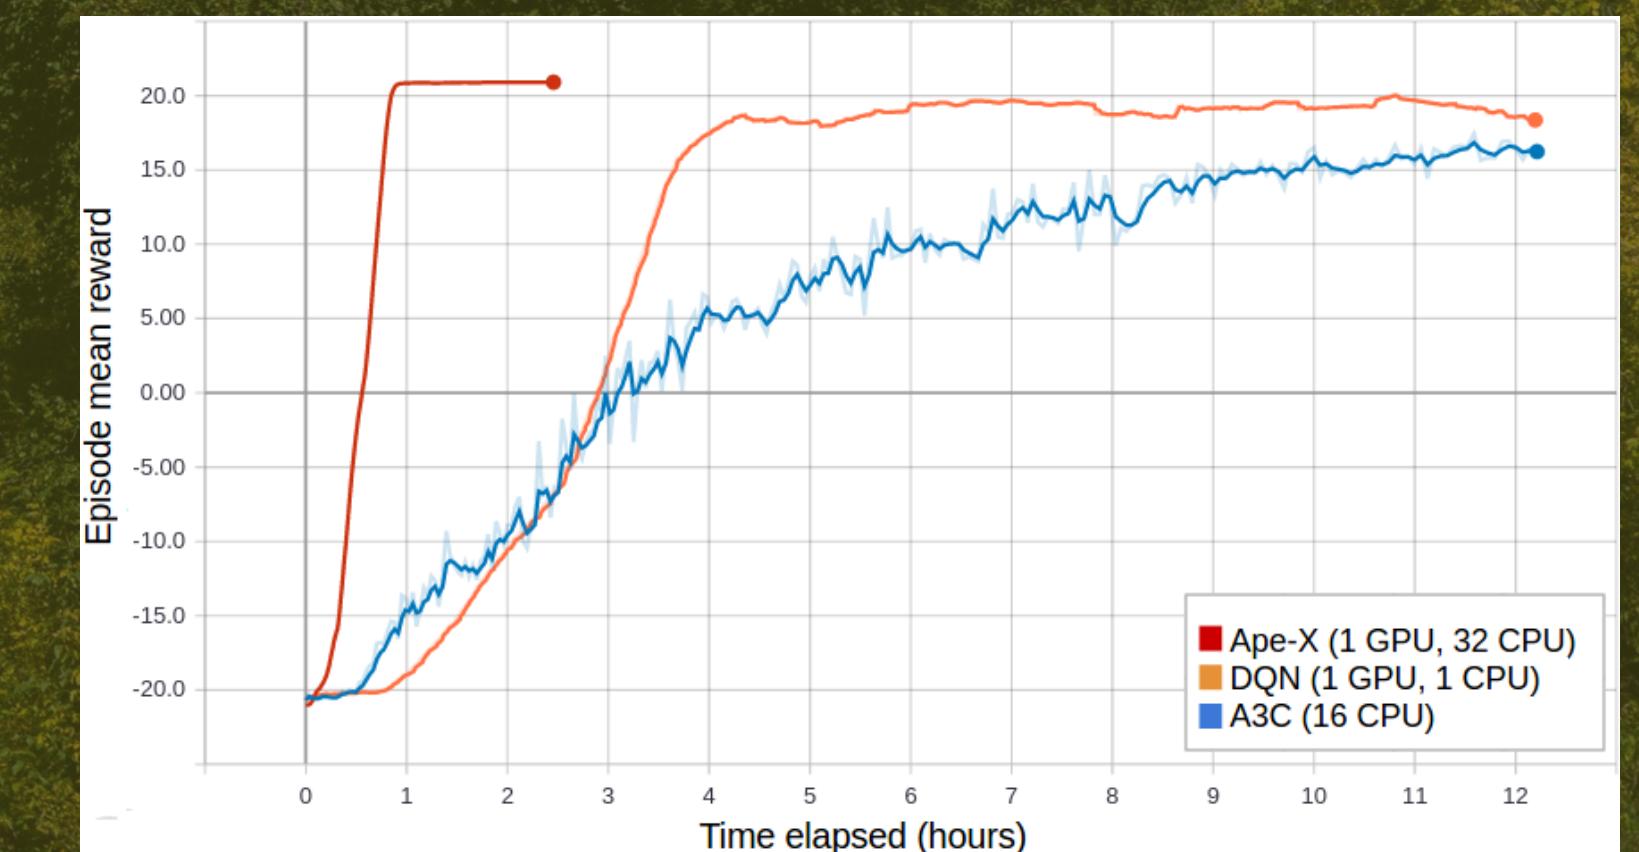
RLLib Provides a Unified Framework for Scalable RL that Doesn't Compromise on Performance

Distributed PPO



Evolution
Strategies

Ape-X Distributed
DQN, DDPG



A Broad Range of Popular Algorithms

- High-throughput architectures
 - Distributed Prioritized Experience Replay (Ape-X)
 - Importance Weighted Actor-Learner Architecture (IMPALA)
 - Asynchronous Proximal Policy Optimization (APPO)
- Gradient-based
 - Soft Actor-Critic (SAC)
 - Advantage Actor-Critic (A₂C, A₃C)
 - Deep Deterministic Policy Gradients (DDPG, TD3)
 - Deep Q Networks (DQN, Rainbow, Parametric DQN)
 - Policy Gradients
 - Proximal Policy Optimization (PPO)
- gradient-free
 - Augmented Random Search (ARS)
 - Evolution Strategies
- Multi-agent specific
 - QMIX Monotonic Value Factorisation (QMIX, VDN, IQN)
- Offline
 - Advantage Re-Weighted Imitation Learning (MARWIL)



Amazon SageMaker RL

Reinforcement learning for every developer and data scientist



Amazon SageMaker RL

End-to-end examples for classic RL and real-world RL applications

Robotics

Industrial Control

HVAC

Autonomous Vehicles

Operations

Finance

Games

NLP

RL Environments to model real-world problems

AWS Simulation Environments

Amazon Sumerian

AWS RoboMaker

Open Source Environments

EnergyPlus

RoboSchool

PyBullet

...

Custom Environments

Bring Your Own

Commercial simulators

MATLAB & Simulink

Open AI Gym

RL Toolkits that provide RL agent algorithm implementations

RL-Coach

DQN

PPO

HER

Rainbow

...

RL-Ray RLLib

APEX

ES

IMPALA

A3C

...

Open AI Baselines

TRPO

GAIL

...

...

TensorFlow

MxNet

PyTorch

Chainer

Training Options

Single Machine / Distributed

Local / Remote simulation

CPU / GPU Hardware

SageMaker supported

Customer BYO

Now in Azure

The screenshot shows a Microsoft Docs page titled "Reinforcement learning (preview) with Azure Machine Learning". The page is part of the "Azure Machine Learning Documentation" under the "Machine Learning" category. The main content area features a large heading and a note about the feature being a preview. A sidebar on the left contains navigation links for "Overview", "Tutorials", and "Samples". The Microsoft logo is visible at the top left of the page.

Microsoft | Docs Documentation Learn Q&A Code Samples

Azure Product documentation ▾ Architecture ▾ Learn Azure ▾ Develop ▾ Resources ▾

Azure / Machine Learning Bookmark

Azure Machine Learning Documentation

✓ Overview

- What is Azure Machine Learning?
- Azure Machine Learning vs Studio (classic)
- Architecture & terms

✓ Tutorials

- > Studio
- > Python SDK
- > R SDK
- > Machine Learning CLI
- > Visual Studio Code

> Samples

> Concepts

Reinforcement learning (preview) with Azure Machine Learning

05/05/2020 • 11 minutes to read • 

APPLIES TO:  Basic edition  Enterprise edition [\(Upgrade to Enterprise edition\)](#)

Note

Azure Machine Learning Reinforcement Learning is currently a preview feature. Only Ray and RLLib frameworks are supported at this time.

In this article, you learn how to train a reinforcement learning (RL) agent to play the video game Pong. You will use the open-source Python library [Ray RLlib](#) with Azure Machine Learning to manage the complexity of distributed RL jobs.

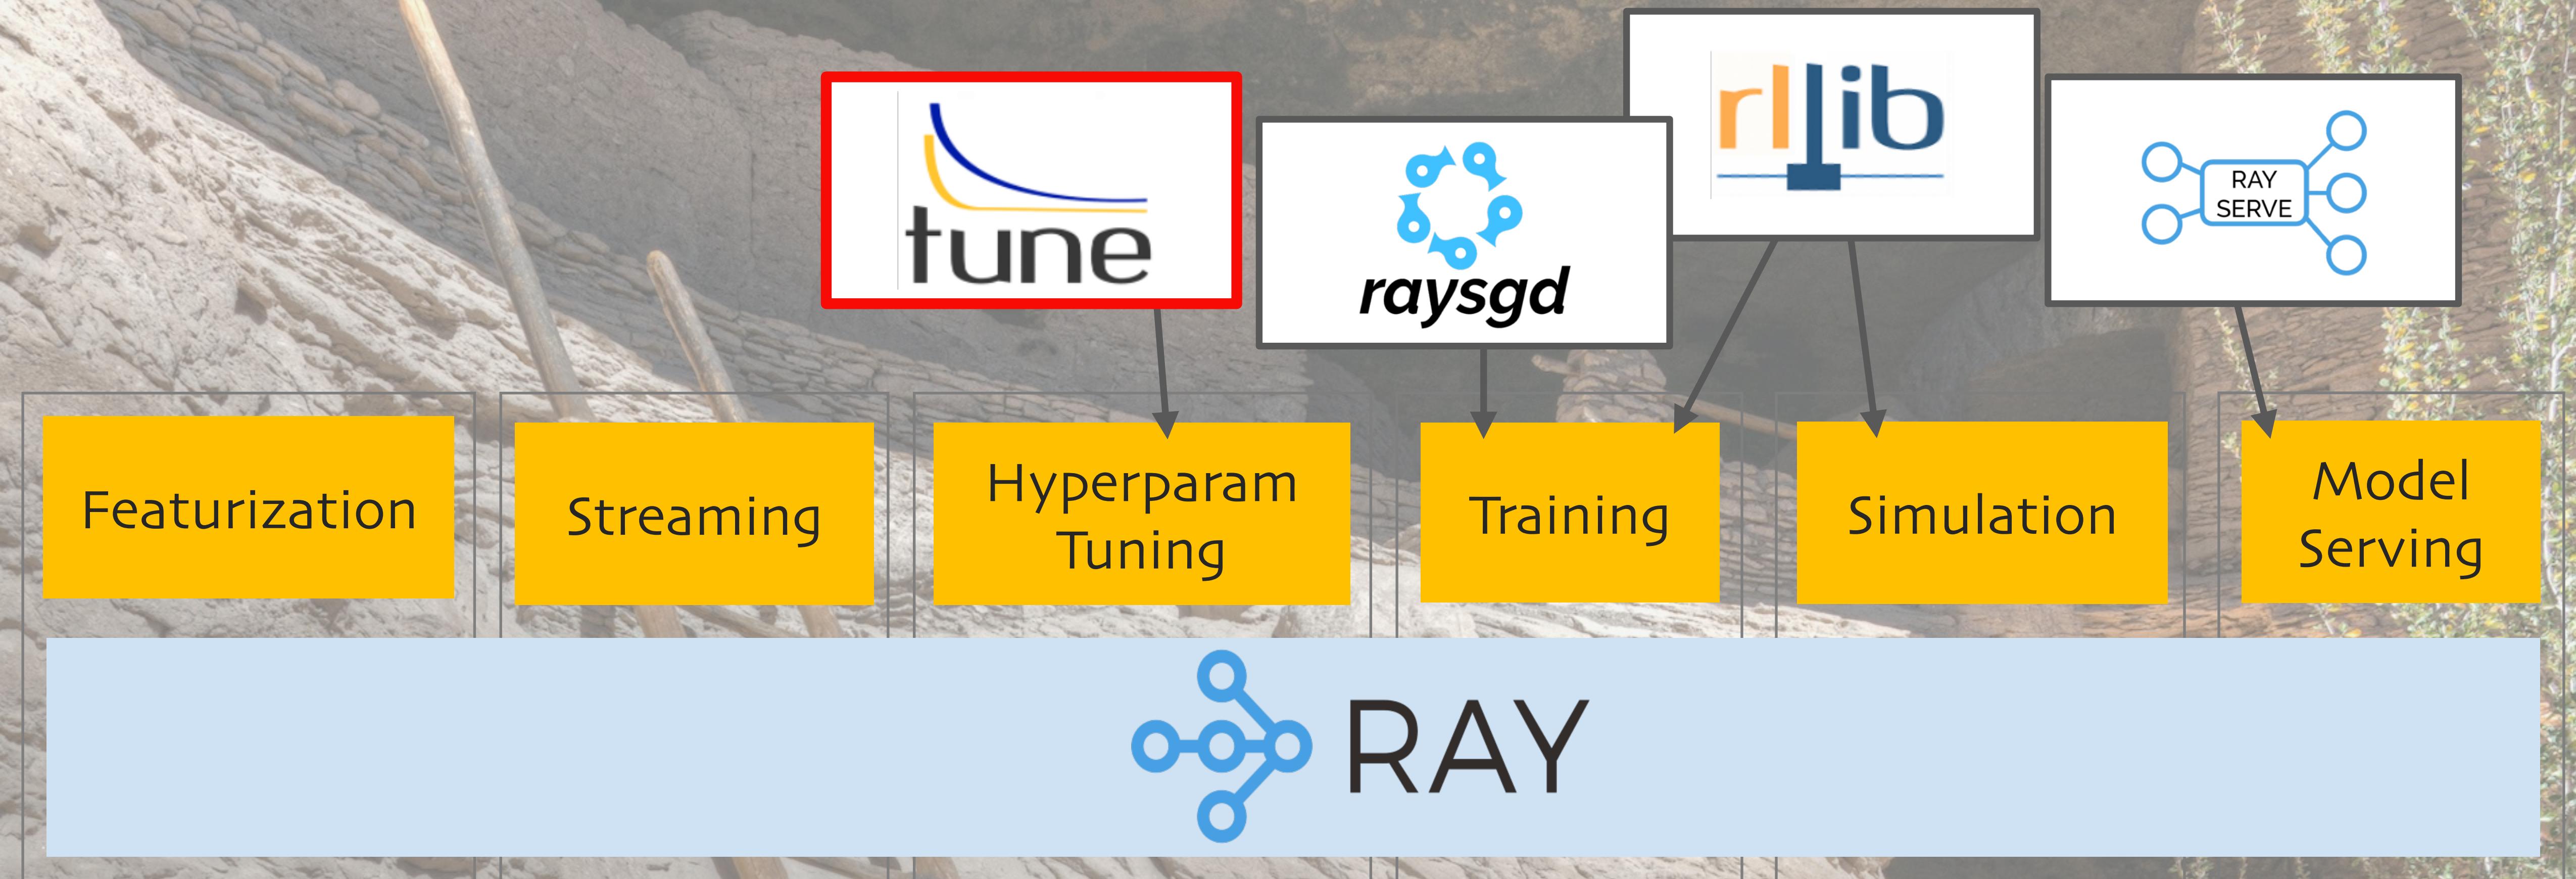
In this article you will learn how to:





@deanwampler

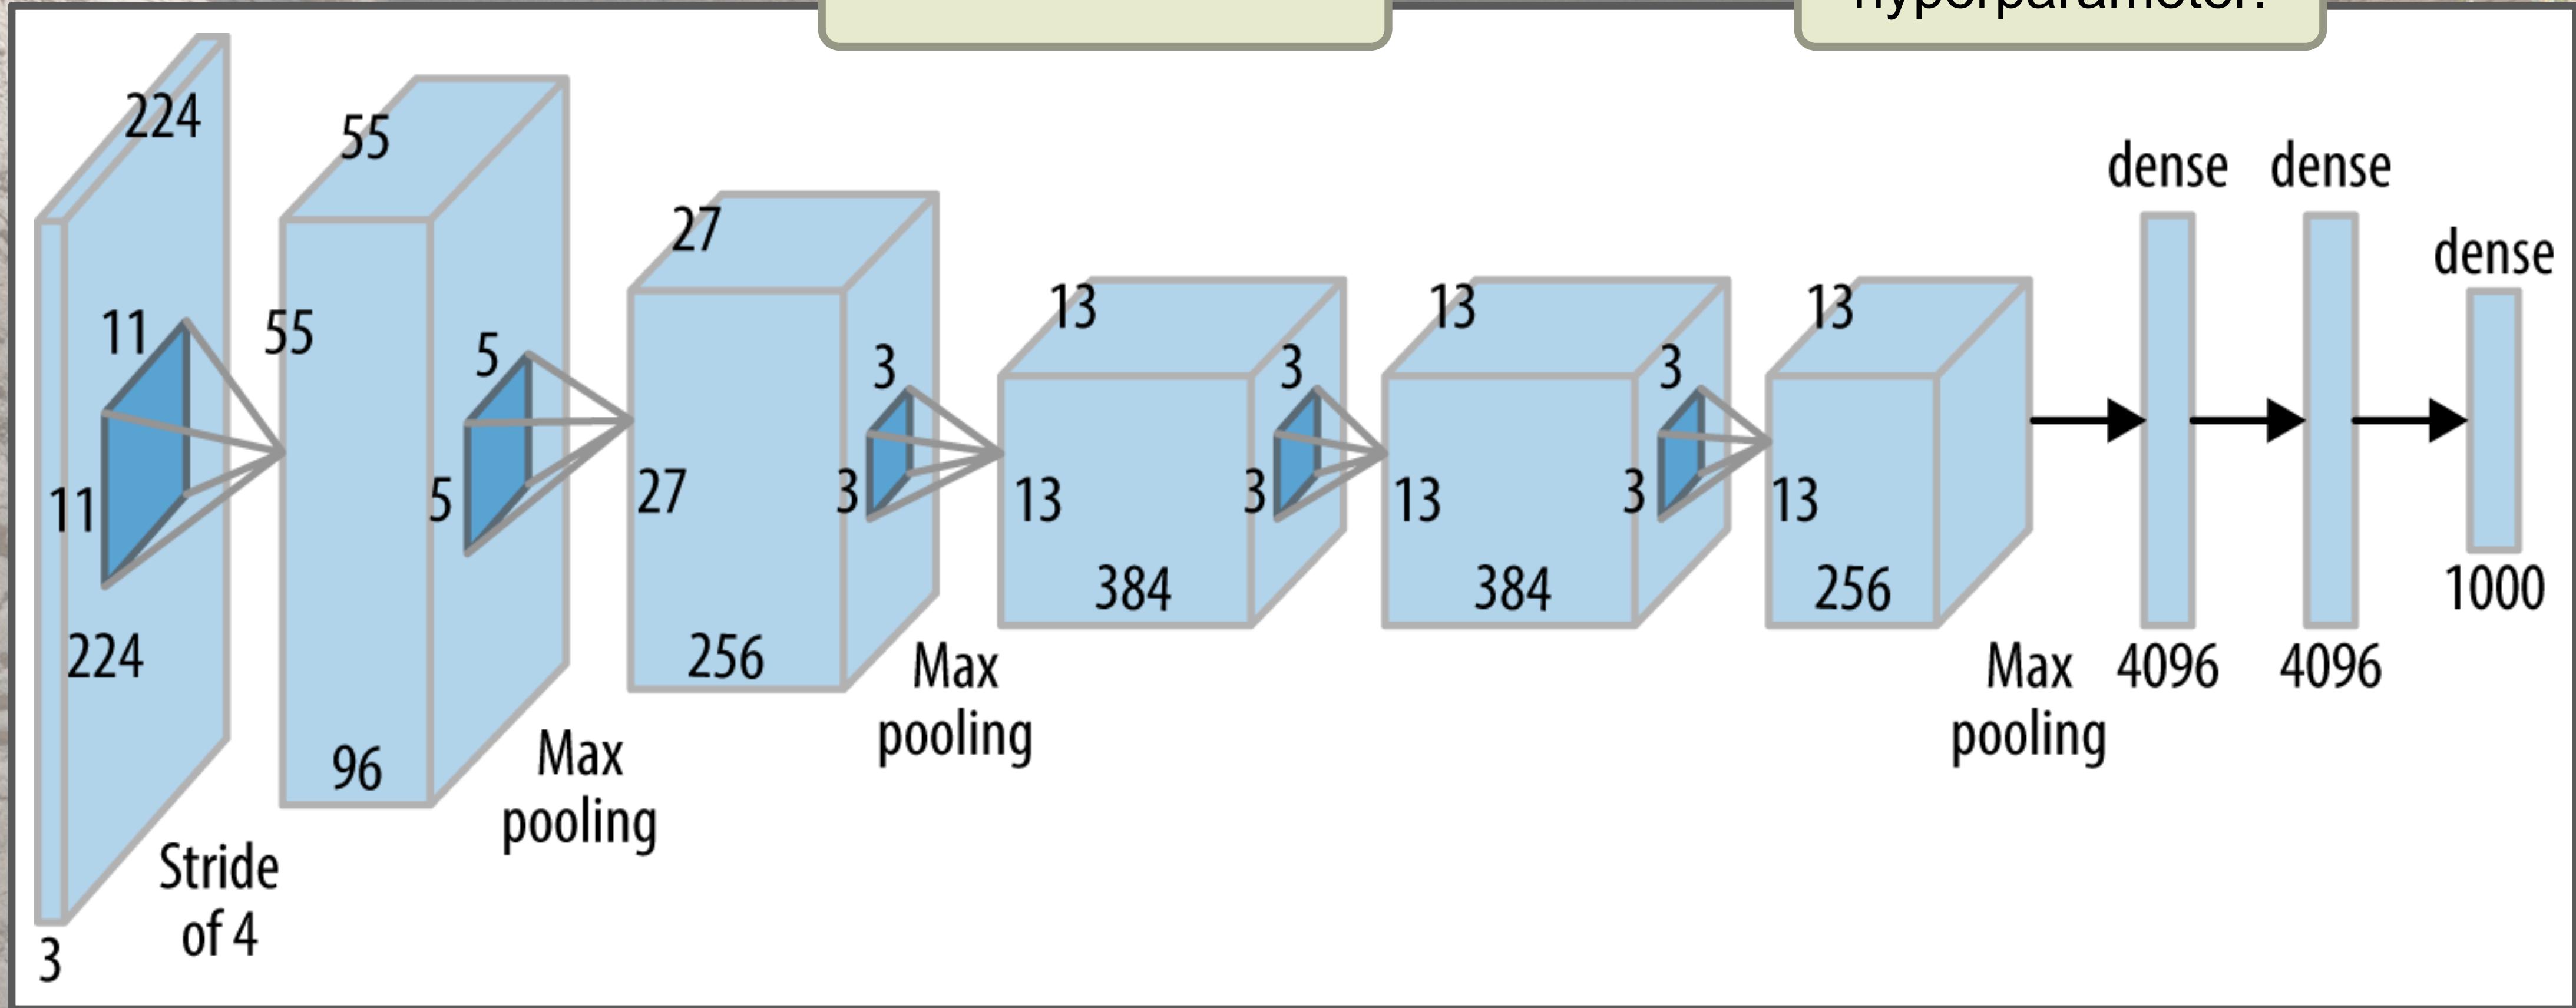
Hyperparameter Tuning - Ray Tune



Nontrivial Example - Neural Networks

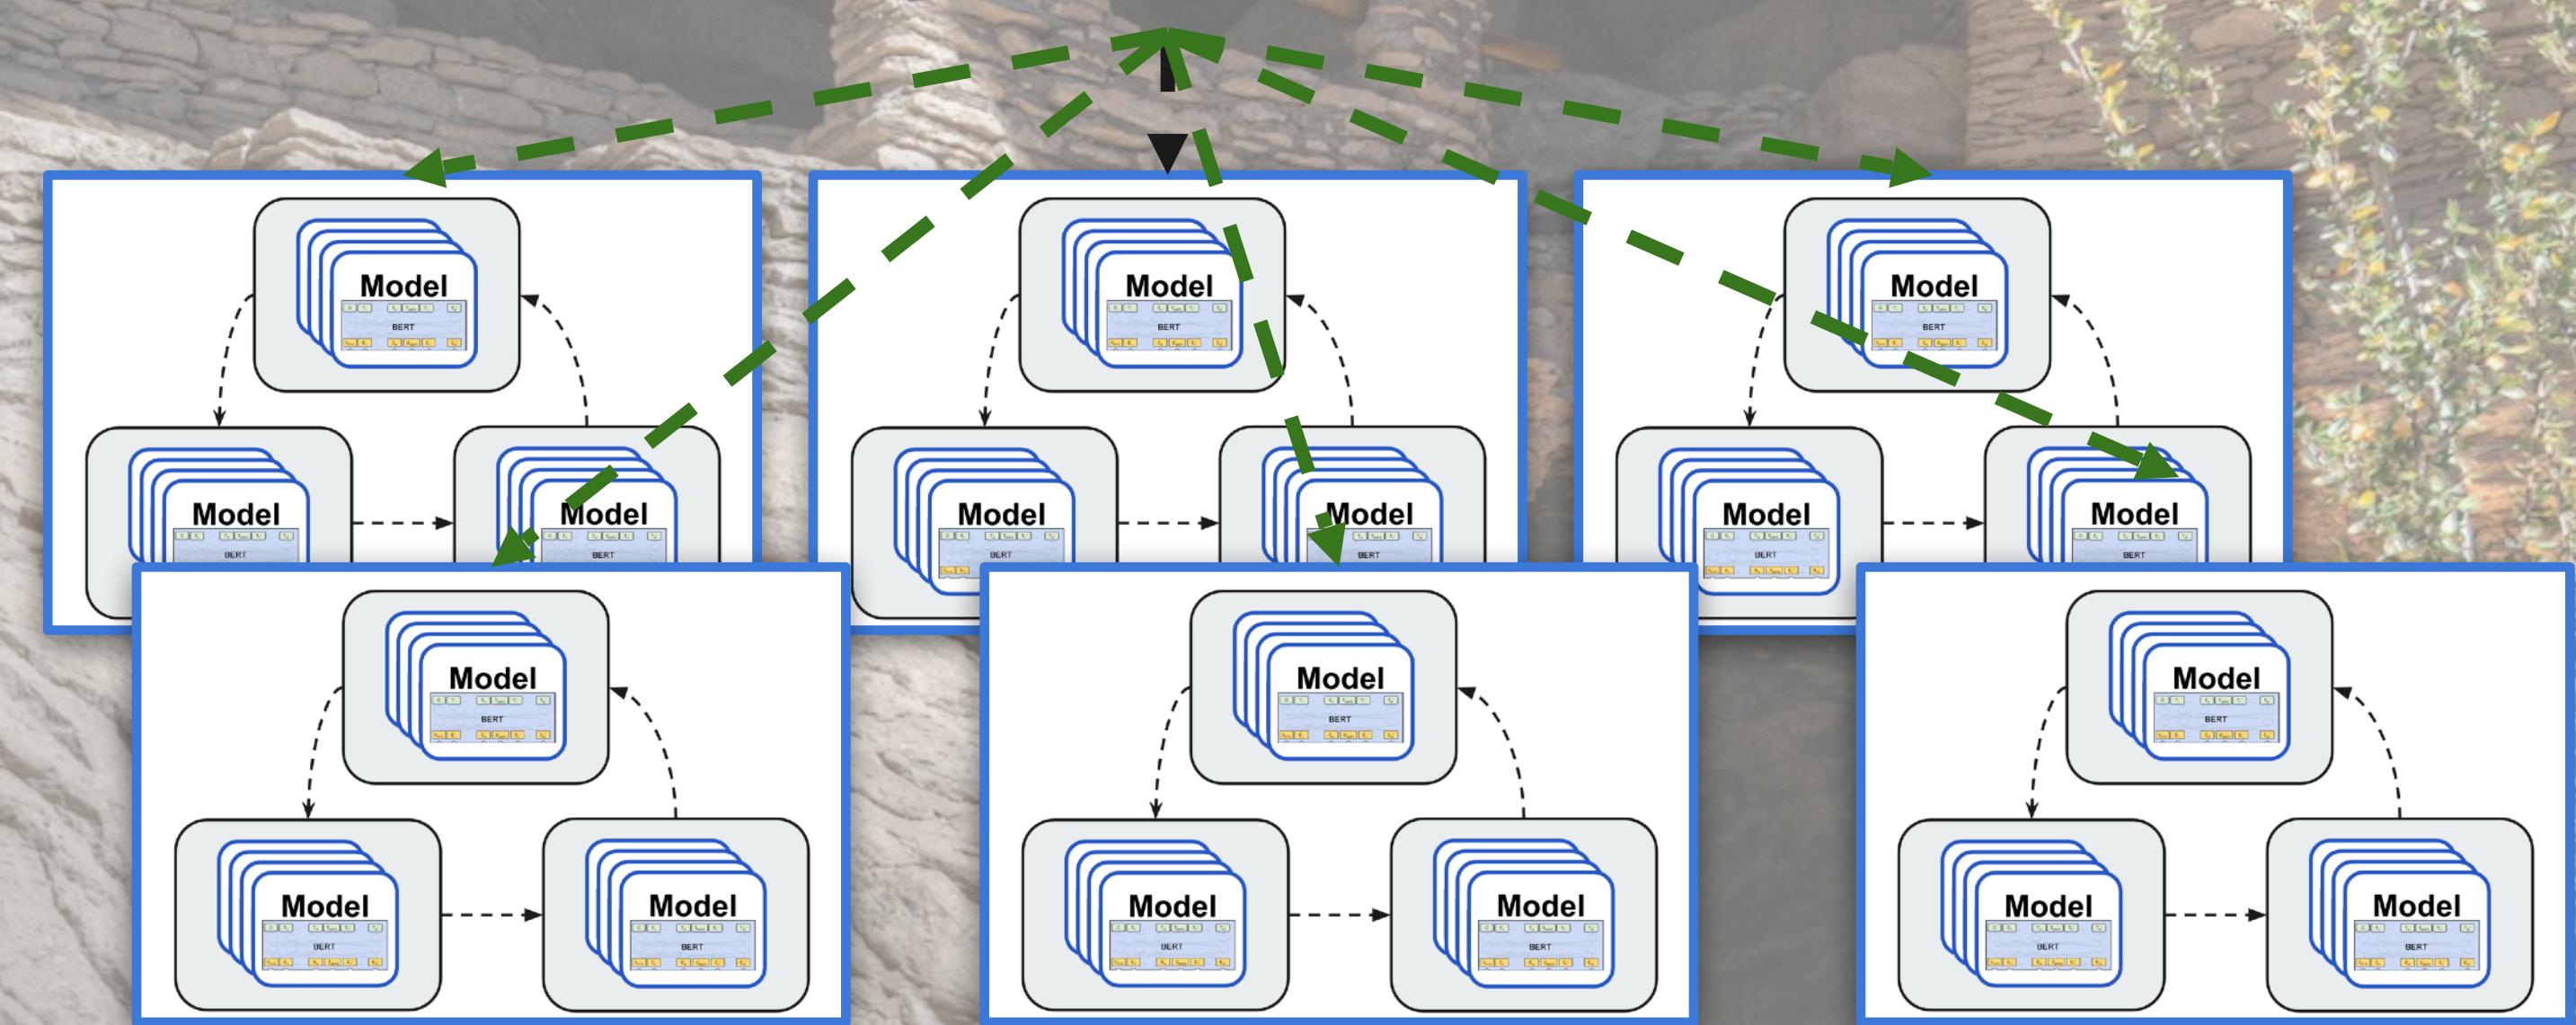
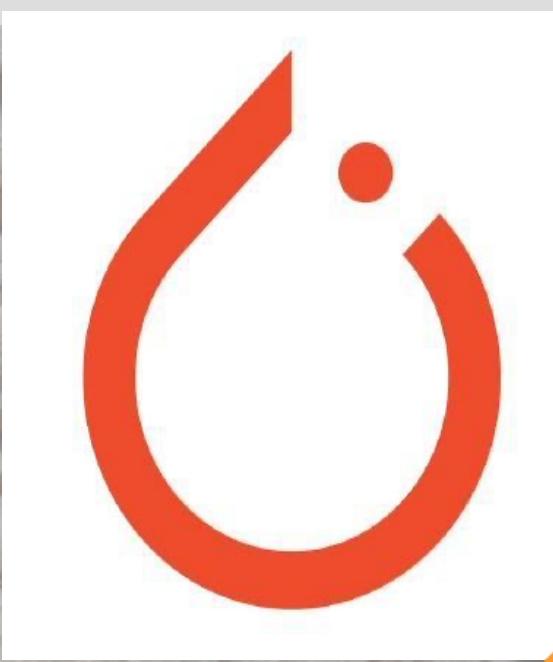
How many layers?
What kinds of layers?

Every number
shown is a
hyperparameter!



Tuning + Distributed Training

```
tune.run(PytorchTrainable,  
 config={  
     "model_creator": PretrainBERT,  
     "data_creator": create_data_loader,  
     "use_gpu": True,  
     "num_replicas": 8,  
     "lr": tune.uniform(0.001, 0.1)  
 },  
 num_samples=100,  
 search_alg=BayesianOptimization()
```





Next Steps: Adopting Ray



@deanwampler

Ray Community and Resources

- ray.io
- Ray Slack: ray-distributed.slack.com
- Tutorials: anyscale.com/academy



Thank You

ray.io

dean@deanwampler.com

[@deanwampler](https://twitter.com/deanwampler)

polyglotprogramming.com/talks

raysummit.org
anyscale.com/events



September 30–October 1