

Reactive Systems: The Why and the How

Dean Wampler, Ph.D.
Typesafe



©Typesafe 2014-2015, All Rights Reserved

O'REILLY®

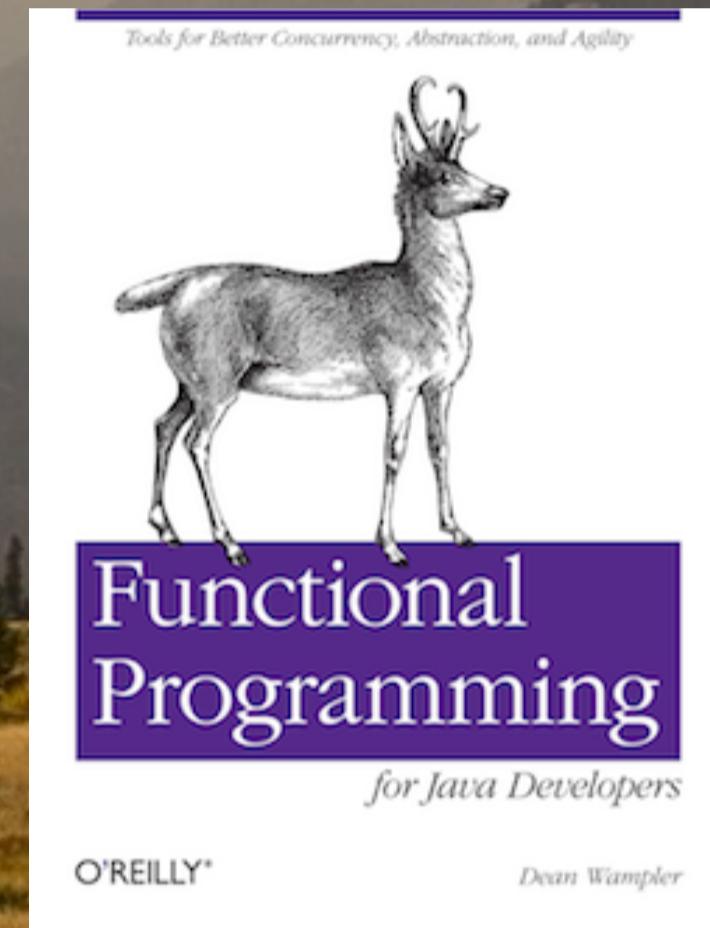
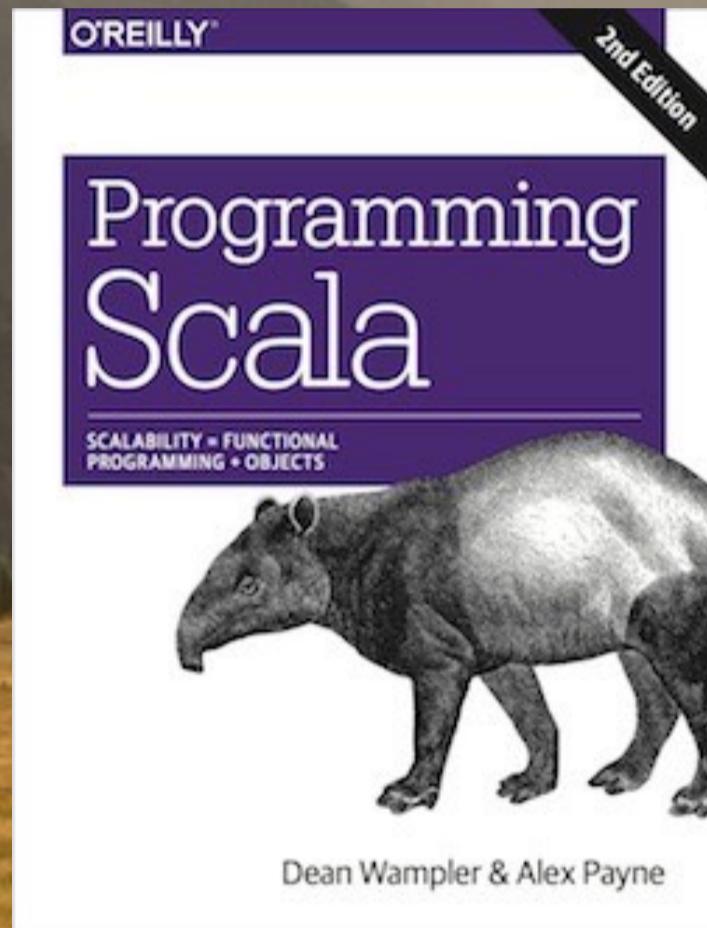
Software Architecture
CONFERENCE

MARCH 16-19, 2015
BOSTON, MA

Tuesday, March 17, 15

Photos from Colorado, Sept. 2014.
All photos are copyright (C) 2000–2015, Dean Wampler, except where otherwise noted. All Rights Reserved.

dean.wampler@typesafe.com
polyglotprogramming.com/talks
@deanwampler





Typesafe Reactive Big Data

typesafe.com/reactive-big-data

3

Tuesday, March 17, 15

This is my role. We're now rolling out commercial support for Spark in non-Hadoop environments. We have other projects in the works. Talk to me if you're interested in what we're doing.

Later today:

Error Handling in Reactive Systems

3:30pm–5:00pm Thursday, 03/19/2015

Reactive and its variants

Location: 304

Tags: reactive



Dean Wampler (Typesafe)

The Reactive Manifesto's "Resilient" trait requires a system to stay responsive when failures happen. I'll discuss how real-world systems do this, starting with in-process techniques in Go, Clojure's core.async, and Reactive Extensions. Next, I'll discuss how some tools prevent common failures in the first place. I'll finish with the Actor model's strategic use of supervisor hierarchies. [Read more](#).

31

Tuesday, March 17, 15

4

This talk is a general overview of Reactive concepts. I'm going to dive into the "Resilient" trait in more detail later today.

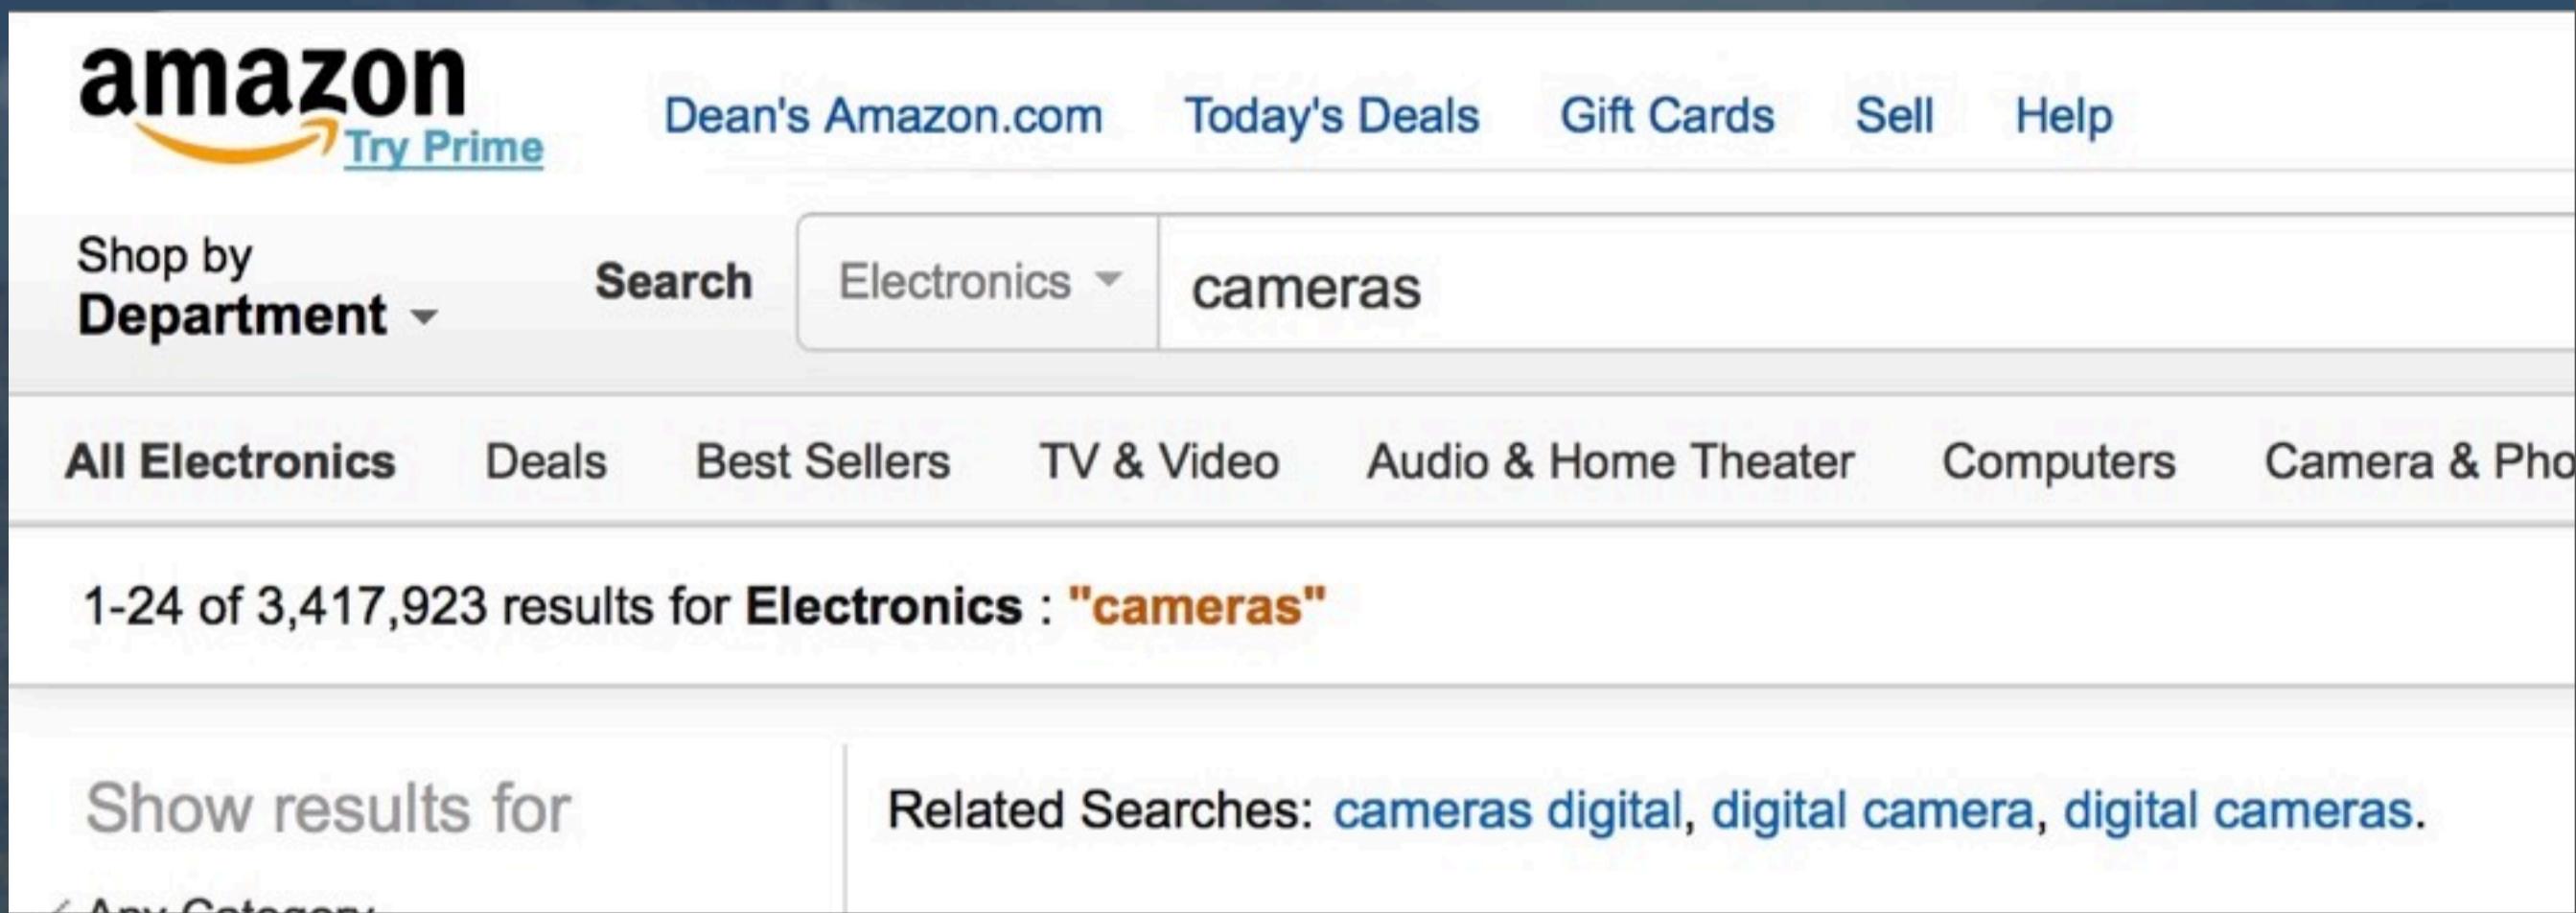
Motivation: eCommerce

5

Tuesday, March 17, 15

Let's motivate the notion of "reactive" systems by exploring some common scenarios we see today.

Cyber Monday?

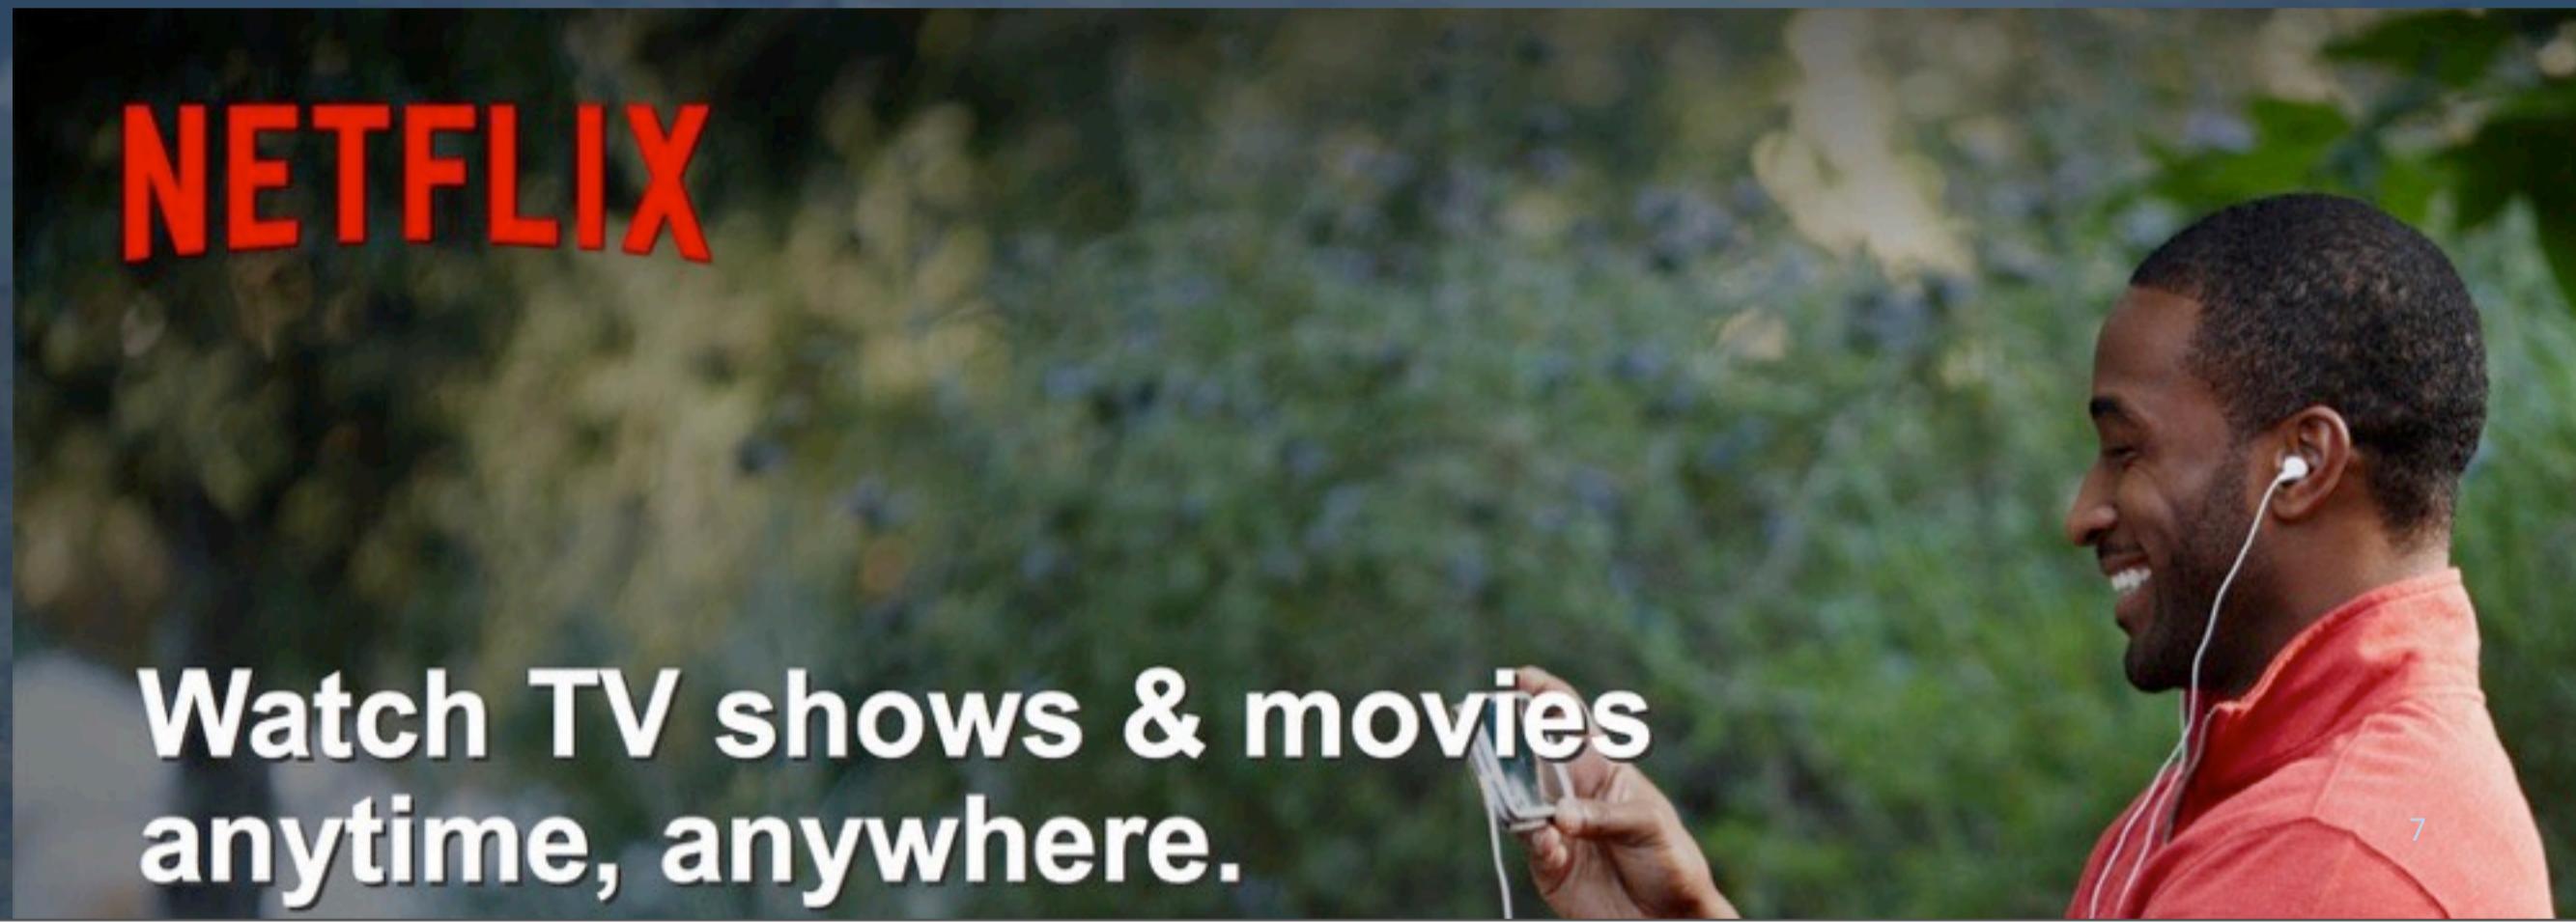


A screenshot of the Amazon website homepage. The header features the Amazon logo with "Try Prime" underneath. To the right are links for "Dean's Amazon.com", "Today's Deals", "Gift Cards", "Sell", and "Help". Below the header is a search bar with "Shop by Department" dropdown, a "Search" button, a "Electronics" dropdown set to "cameras", and a main search input field containing "cameras". Underneath the search bar are category links: "All Electronics", "Deals", "Best Sellers", "TV & Video", "Audio & Home Theater", "Computers", and "Camera & Photo". A message below the search bar states "1-24 of 3,417,923 results for Electronics : cameras". On the left, there's a "Show results for" link and a "Any Category" link. On the right, related searches include "cameras digital", "digital camera", and "digital cameras".

Tuesday, March 17, 15

Your online store needs to scale up and down with demand. It needs to degrade gracefully if some service components are lost or disconnected by a network partition. For example, if the canonical catalog “disappears” behind a network partition, it’s probably better to sell off a stale copy locally.
photo: Amazon home page, <http://amazon.com>.

On demand?



Tuesday, March 17, 15

Netflix has extreme scale challenges, but they have become the industry leader in building highly resilient, scalable services. What happens when a new season of “House of Cards” is released?
photo: Netflix home page, <http://netflix.com>.



Motivation: Internet of Things

8

Tuesday, March 17, 15

Internet of Things has several categories of applications, each of which has needs that motivate reactive programming.

Medical Devices, IT Systems

Phone home:

- Upload data
 - Usage patterns
 - Predictive diagnostics
- Fetch patches



9

Tuesday, March 17, 15

Medical devices upload test results (e.g., ultrasound images and video) to servers. Med. devices and IT systems send requests for automated updates, send the equivalent of “click-stream” data used to assess usability, etc., and increasingly send metrics used to predict potential HW failures or other service needs.

ultrasound photo: <http://www.usa.philips.com/healthcare-product/HC795054/hd11-xe-ultrasound-system>

switch photo: <http://networklessons.com/switching/introduction-to-vtp-vlan-trunking-protocol/>

Medical Devices, IT Systems

Characteristics:

- Stable to intermittent network connectivity
- One way and two way
- Mixed bandwidth



10

Tuesday, March 17, 15

A mobile scanner might move in and out of WiFi zones, so caching data is necessary. IT appliances are (hopefully) always online. Some data is one way, like diagnostic info for predictive analytics, while data uploads or patch requests need acknowledgement. Bandwidth can vary.

Aircraft Engines

Phone home:

- Upload telemetry
 - Predictive diagnostics
 - Redundant tracking data!



11

Tuesday, March 17, 15

Each engine collects 0.5–1TB of data per flight. Most is currently tossed! Our only clue about the final resting place of Malaysian Airlines Flight 370 was engine telemetry picked up by satellites and used to analyze possible routes.

Trucks, Farm Equipment

GPS Tracking:

- Optimize routing, fuel use, etc.
- Spy on drivers?
- Per plant tracking!



Tuesday, March 17, 15

Track data to optimize routing, minimize fuel use with shortest path and/or delivering heaviest items first. Ensure drivers are obeying the rules of the road and company policies. Some farm equipment planting, watering, and fertilizing gear now tracks data per plant!

UPS truck photo: http://en.wikipedia.org/wiki/United_Parcel_Service

Planter/seeder photo: http://www.deere.com/en_US/products/equipment/planting_and_seeding_equipment/planting_and_seeding_equipment.page

Trucks, Farm Equipment

Connectivity:

- Always along roads
- Intermittent on farms (WiFi in barns?)



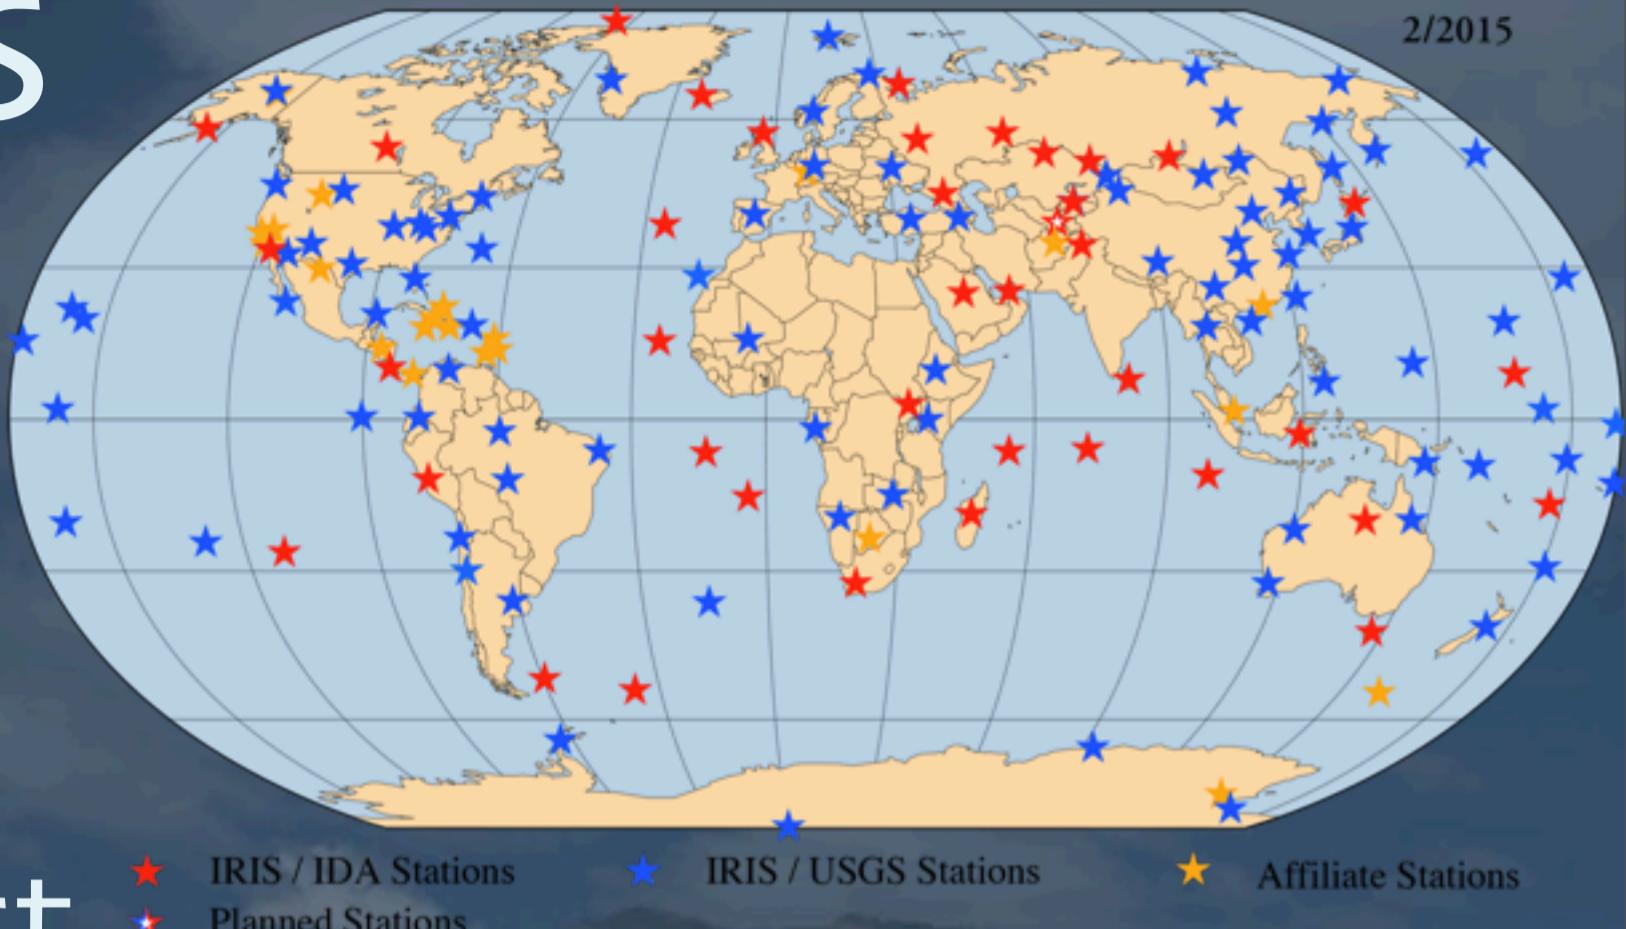
Tuesday, March 17, 15

Some rural areas don't have sufficient wireless data coverage for farm equipment to remain online full time.

Remote Sensors

Human to Real-time Responsiveness:

- Earthquake, nuclear test sensor networks.
- Climate change monitoring



14

Tuesday, March 17, 15

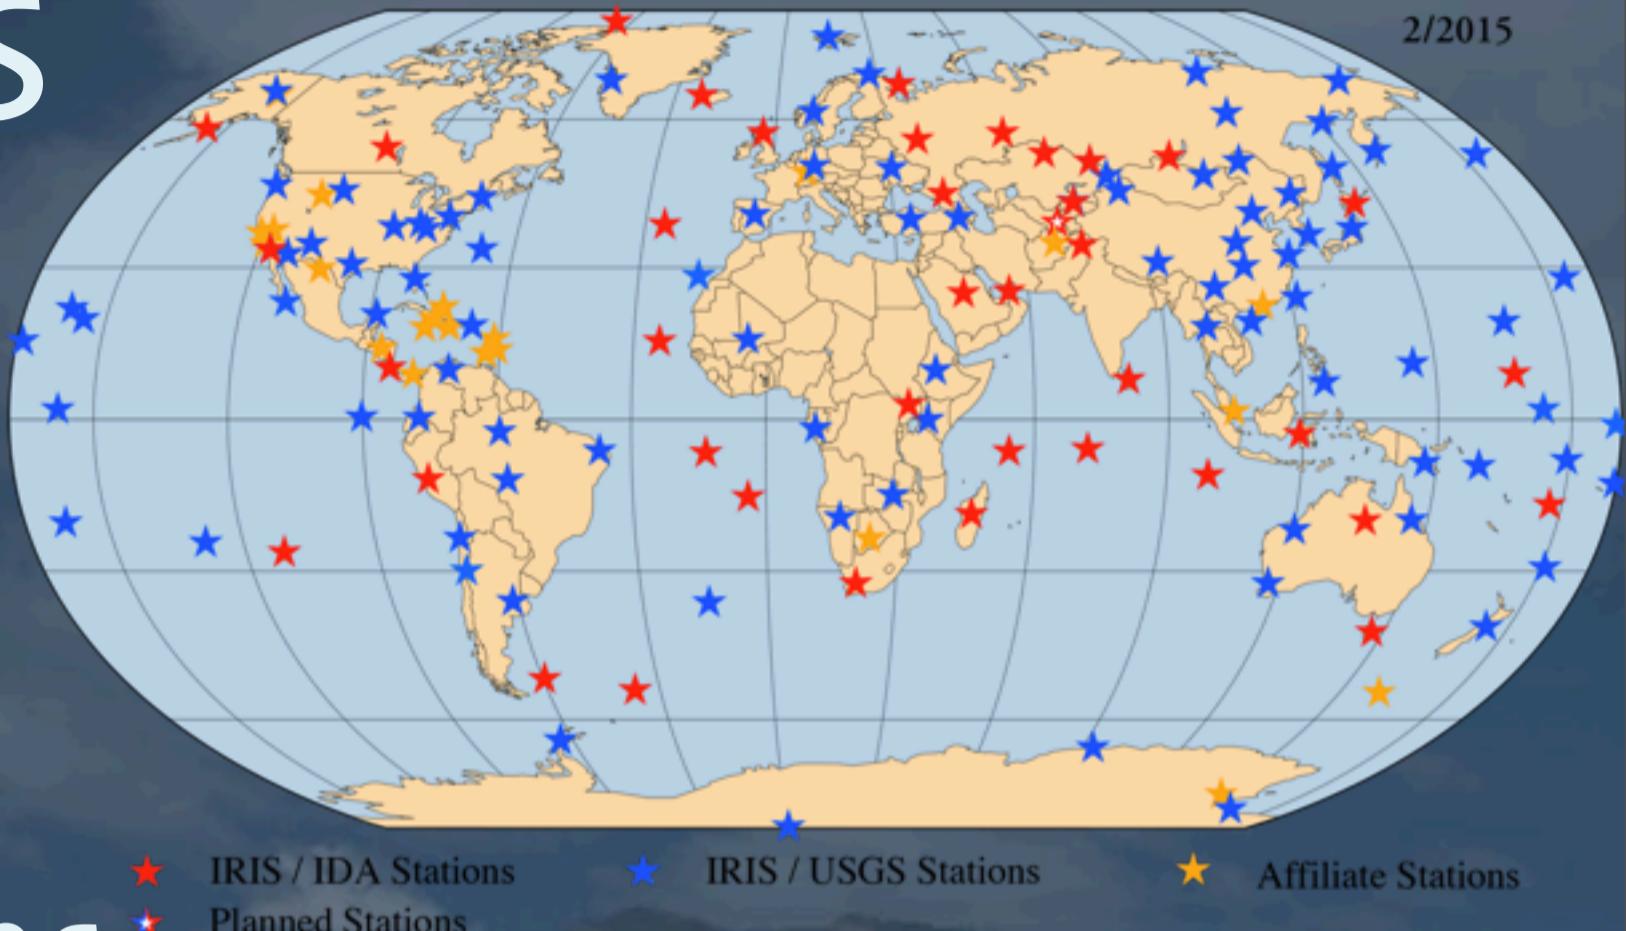
For earthquake sensor networks, you want to get the information to emergency services and community alarm systems in milliseconds. This requires redundant sensors, always on connectivity, and low-latency connections. The amount of data isn't broad. Some networks, like monitoring rainfall or glaciers for climate change studies, might be offline except for once-per-year downloads onsite!

photo: <http://www.iris.edu/hq/programs/gsn>

Remote Sensors

Characteristics:

- Redundant sensors
- Low-latency connections
- Low-bandwidth requirements



15

Tuesday, March 17, 15

For earthquake sensor networks, you want to get the information to emergency services and community alarm systems in milliseconds. This requires redundant sensors, always on connectivity, and low-latency connections. The amount of data isn't broad. Some networks, like monitoring rainfall or glaciers for climate change studies, might be offline except for once-per-year downloads onsite!

Robotics

Connectivity:

- Two-way, but time of flight matters!
- Autonomous?



Tuesday, March 17, 15

Some rural areas don't have sufficient wireless data coverage for farm equipment to remain online full time.

Quadcopter photo: <http://www.dji.com/product/phantom>

Mars rover photo: http://en.wikipedia.org/wiki/Mars_Exploration_Rover

Health Monitoring

Characteristics:

- Occasional to always-on connectivity
- Detect health emergencies: call for help?



17

Tuesday, March 17, 15

Health monitoring tools are most popular for gathering activity data and some vital signs for analysis later. Some monitors are designed to detect medical emergencies and call for help when needed.

photo: <http://www.fitbit.com/force>

Home Automation

Characteristics:

- Morning packet storms
- Fire & break-in detection: automatic notification of authorities



18

Tuesday, March 17, 15

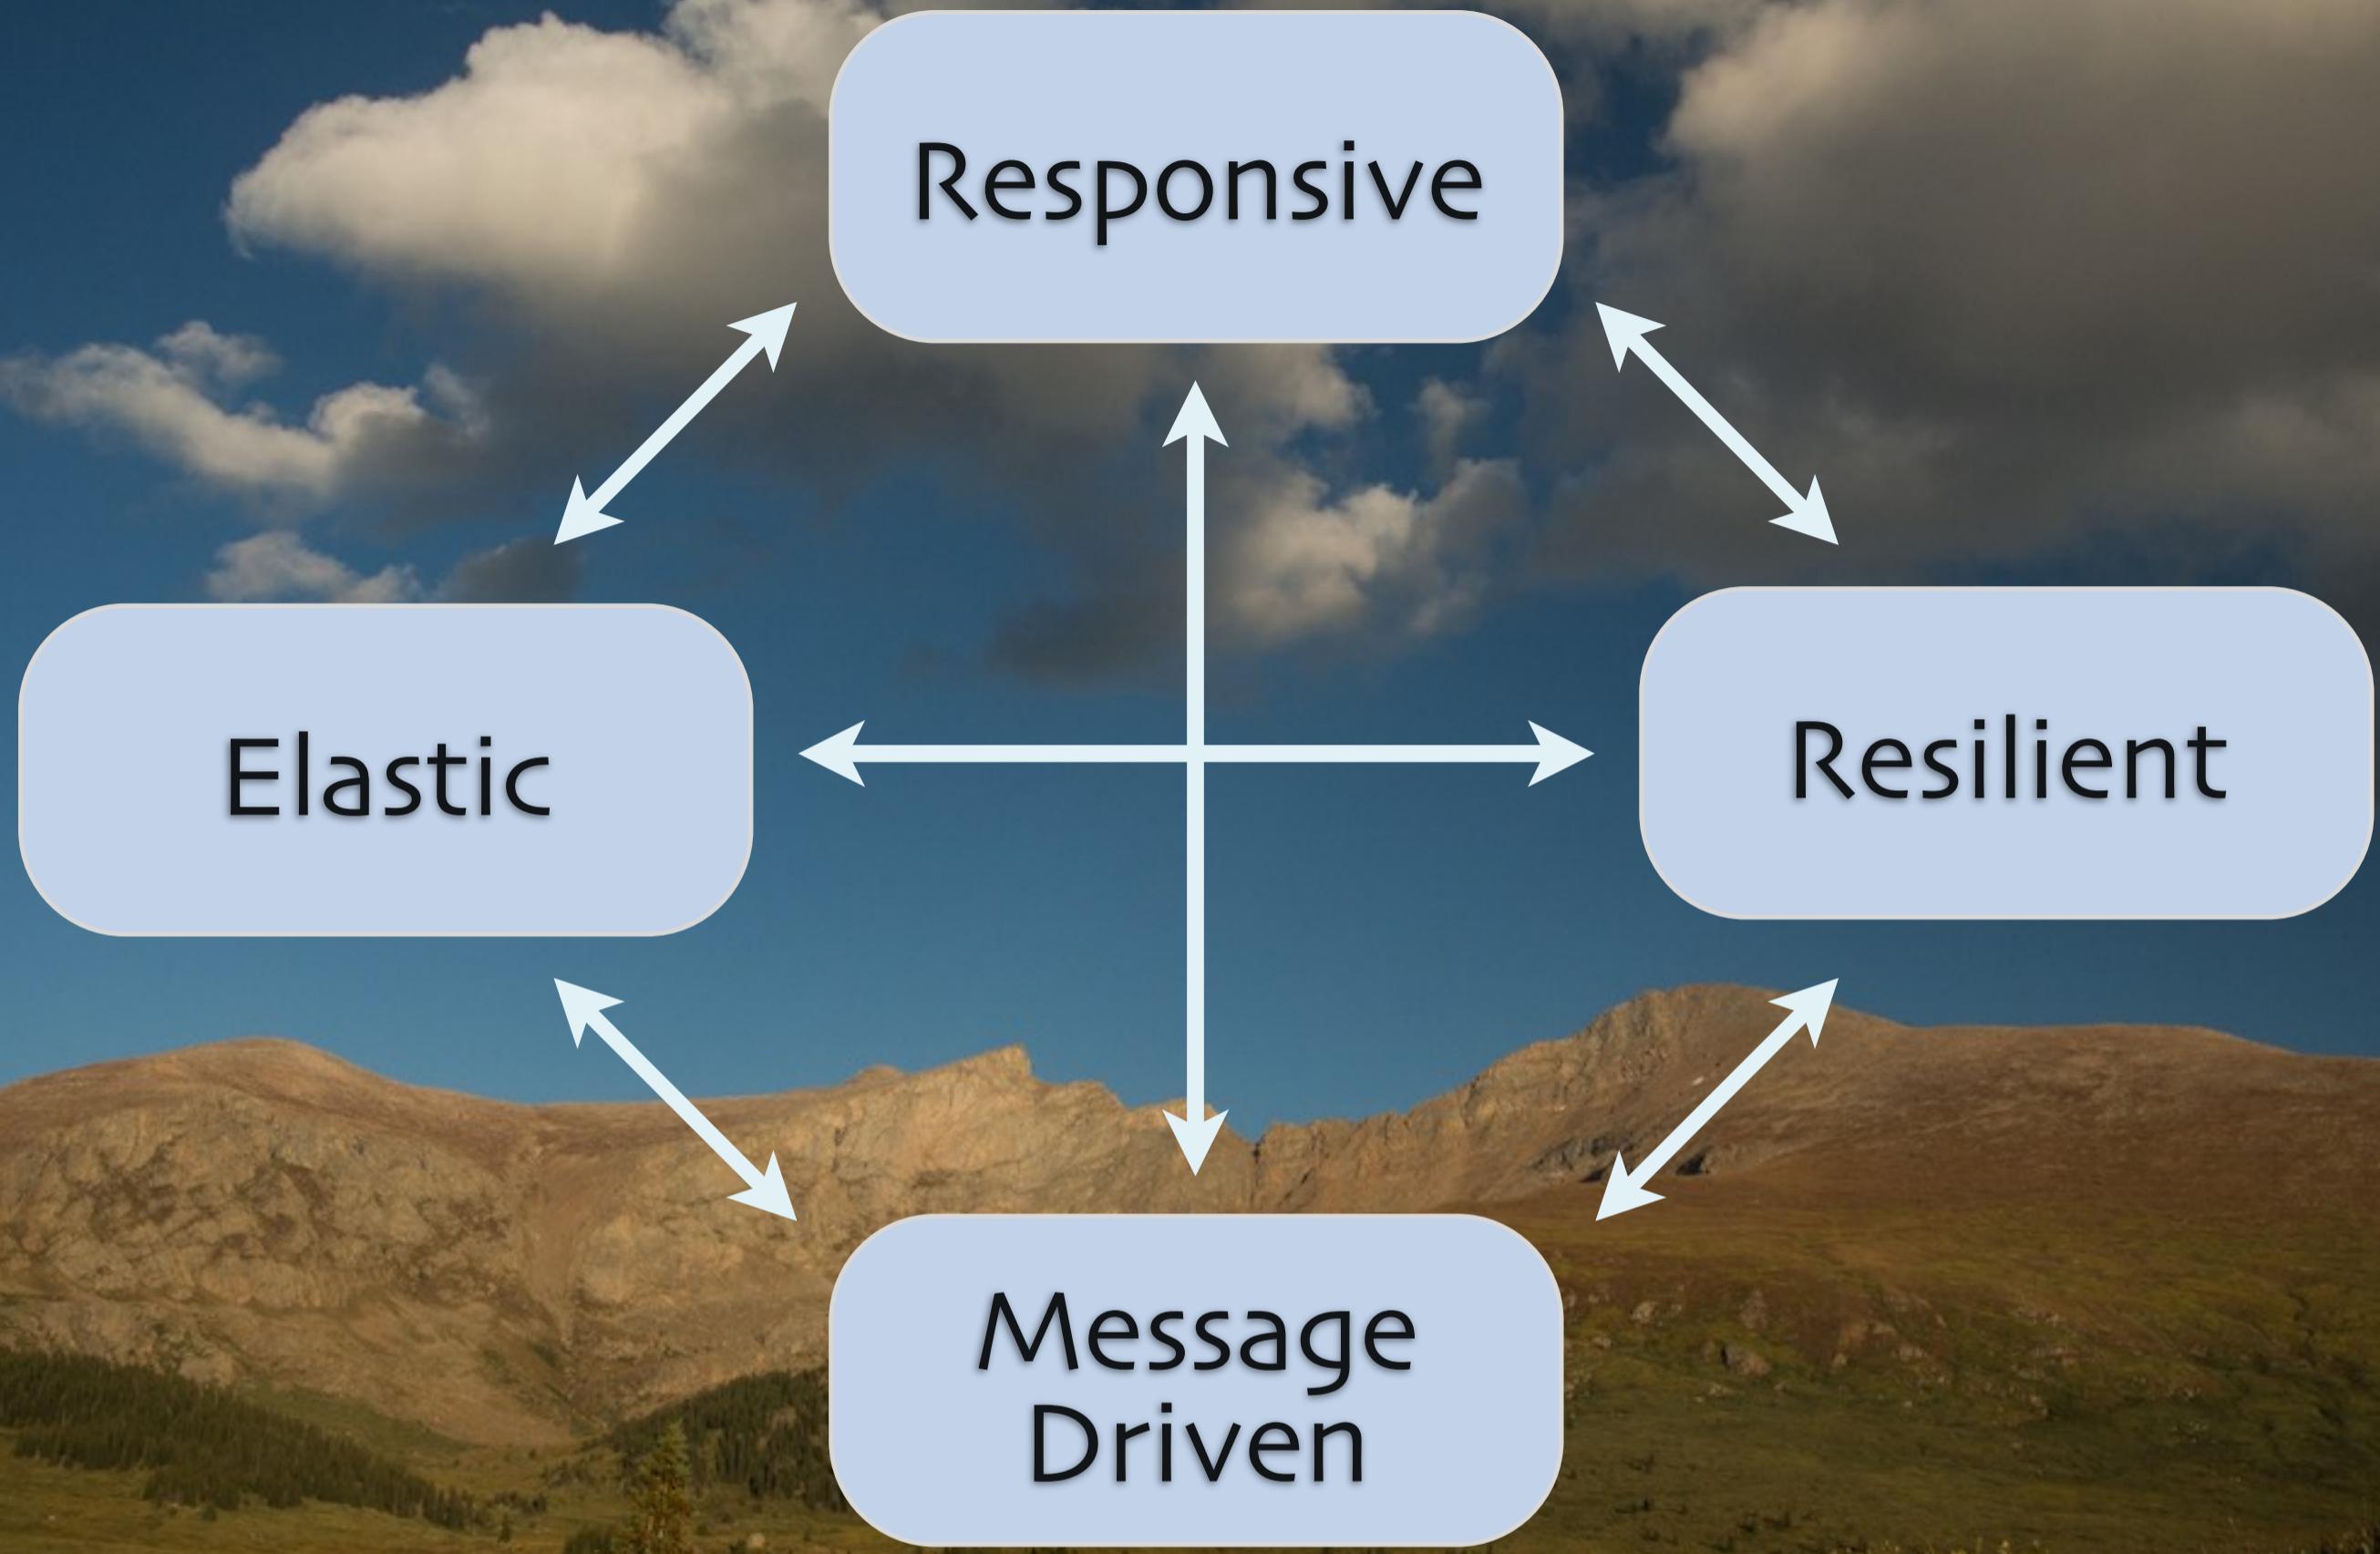
photo: <https://store.nest.com/product/thermostat/>

Reactive Systems

19

Tuesday, March 17, 15

The idea of Reactive Systems emerged to catalog universally common characteristics of the systems we have to build to support these scenarios.



20

Tuesday, March 17, 15

The four characteristics of Reactive Systems... as articulated by the Reactive Manifesto, which attempts to codify lessons learned across many projects, industries, and years building highly available, scalable, and reliable systems.

The Reactive Manifesto

Published on September 16 2014. (v2.0)

Organisations working in disparate domains are independently discovering patterns for building software that look the same. These systems are more robust, more resilient, more flexible and better positioned to meet modern demands.

These changes are happening because application requirements have changed dramatically in recent years. Only a few years ago a large application had tens of servers, seconds of response time, hours of offline maintenance and gigabytes of data. Today applications are deployed on everything from mobile devices to cloud-based clusters running thousands of multi-core processors. Users expect millisecond response times and 100% uptime. Data is measured in Petabytes. Today's demands are simply not met by yesterday's software architectures.

We believe that a coherent approach to systems architecture is needed, and we believe that all necessary aspects are already recognised individually: we want systems that are Responsive, Resilient, Elastic and Message Driven. We call these Reactive Systems.

Systems built as Reactive Systems are more flexible, loosely-coupled and scalable. This makes them easier to develop and amenable to change. They are significantly more tolerant of failure and when failure does occur they meet it with elegance rather than disaster. Reactive Systems are highly responsive, giving users effective interactive feedback.

Tuesday, March 17, 15

The four characteristics of Reactive Systems... as articulated by the Reactive Manifesto, which attempts to codify lessons learned across many projects, industries, and years building highly available, scalable, and reliable systems.



Myths

22

Tuesday, March 17, 15



Myths

This is new.

23

Tuesday, March 17, 15

The RM attempts to codify lessons learned over many years in many scenarios. It's not new. It doesn't claim to be new.



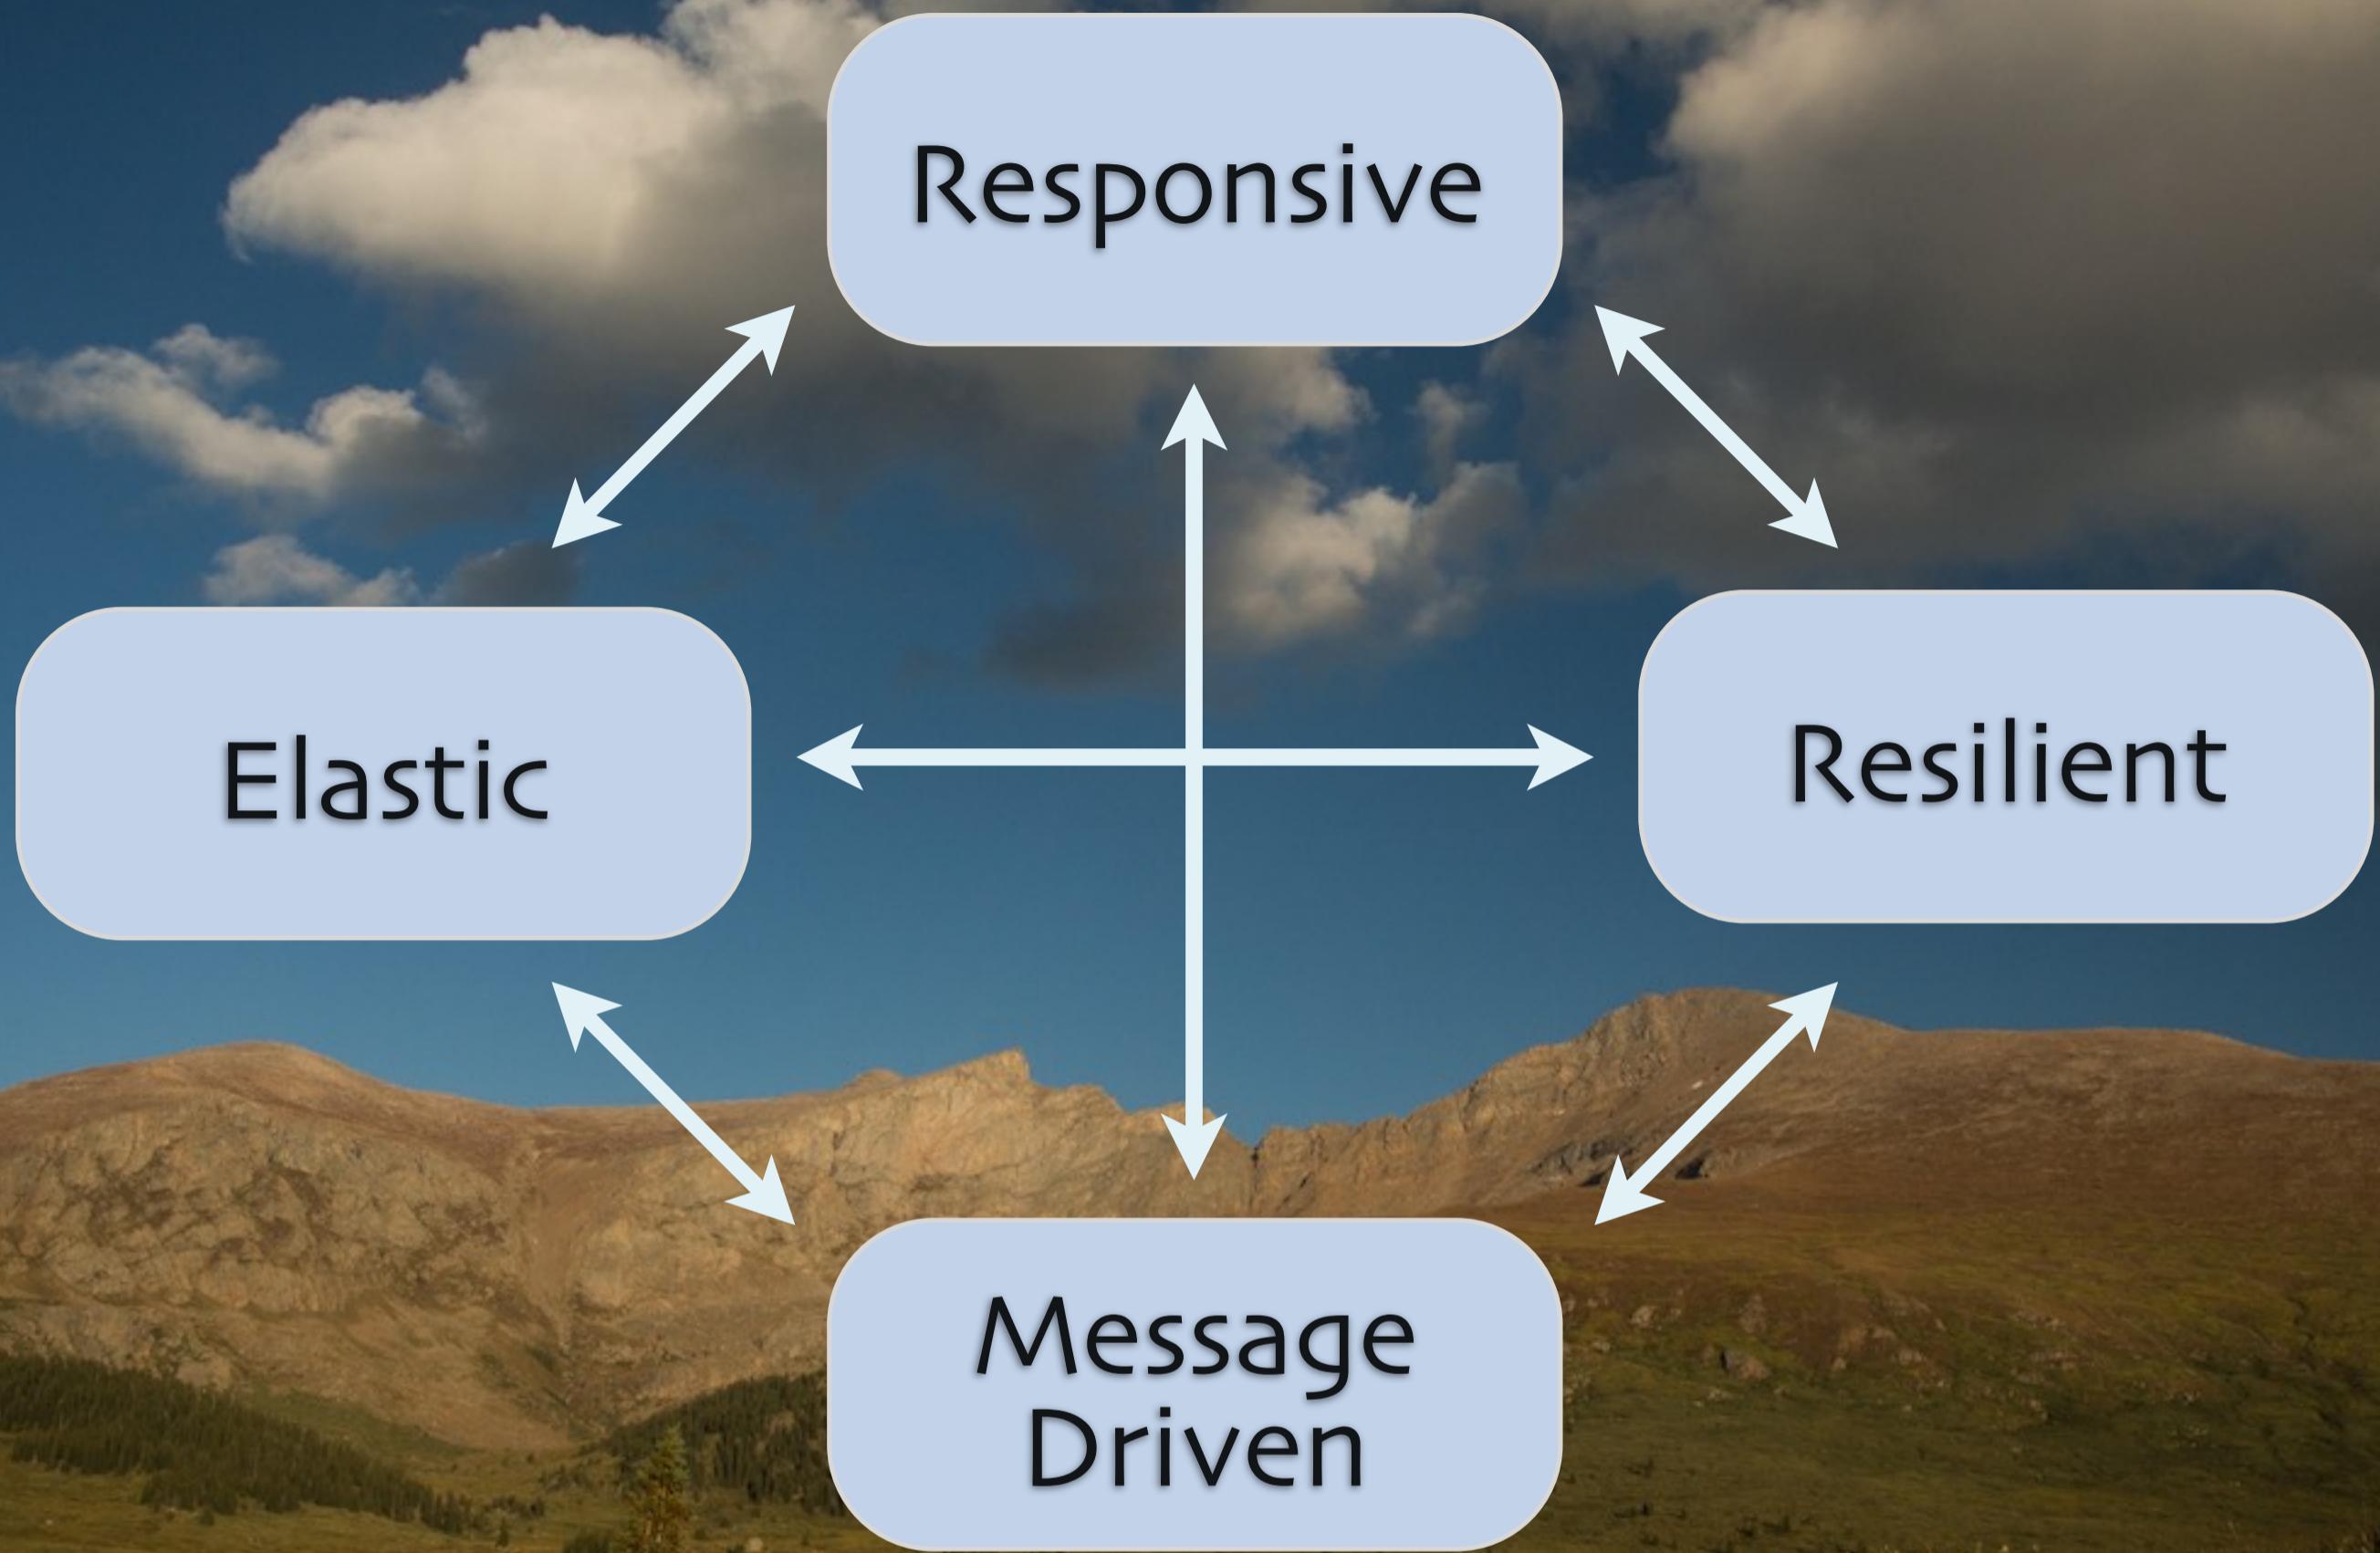
Myths

This is Typesafe marketing.

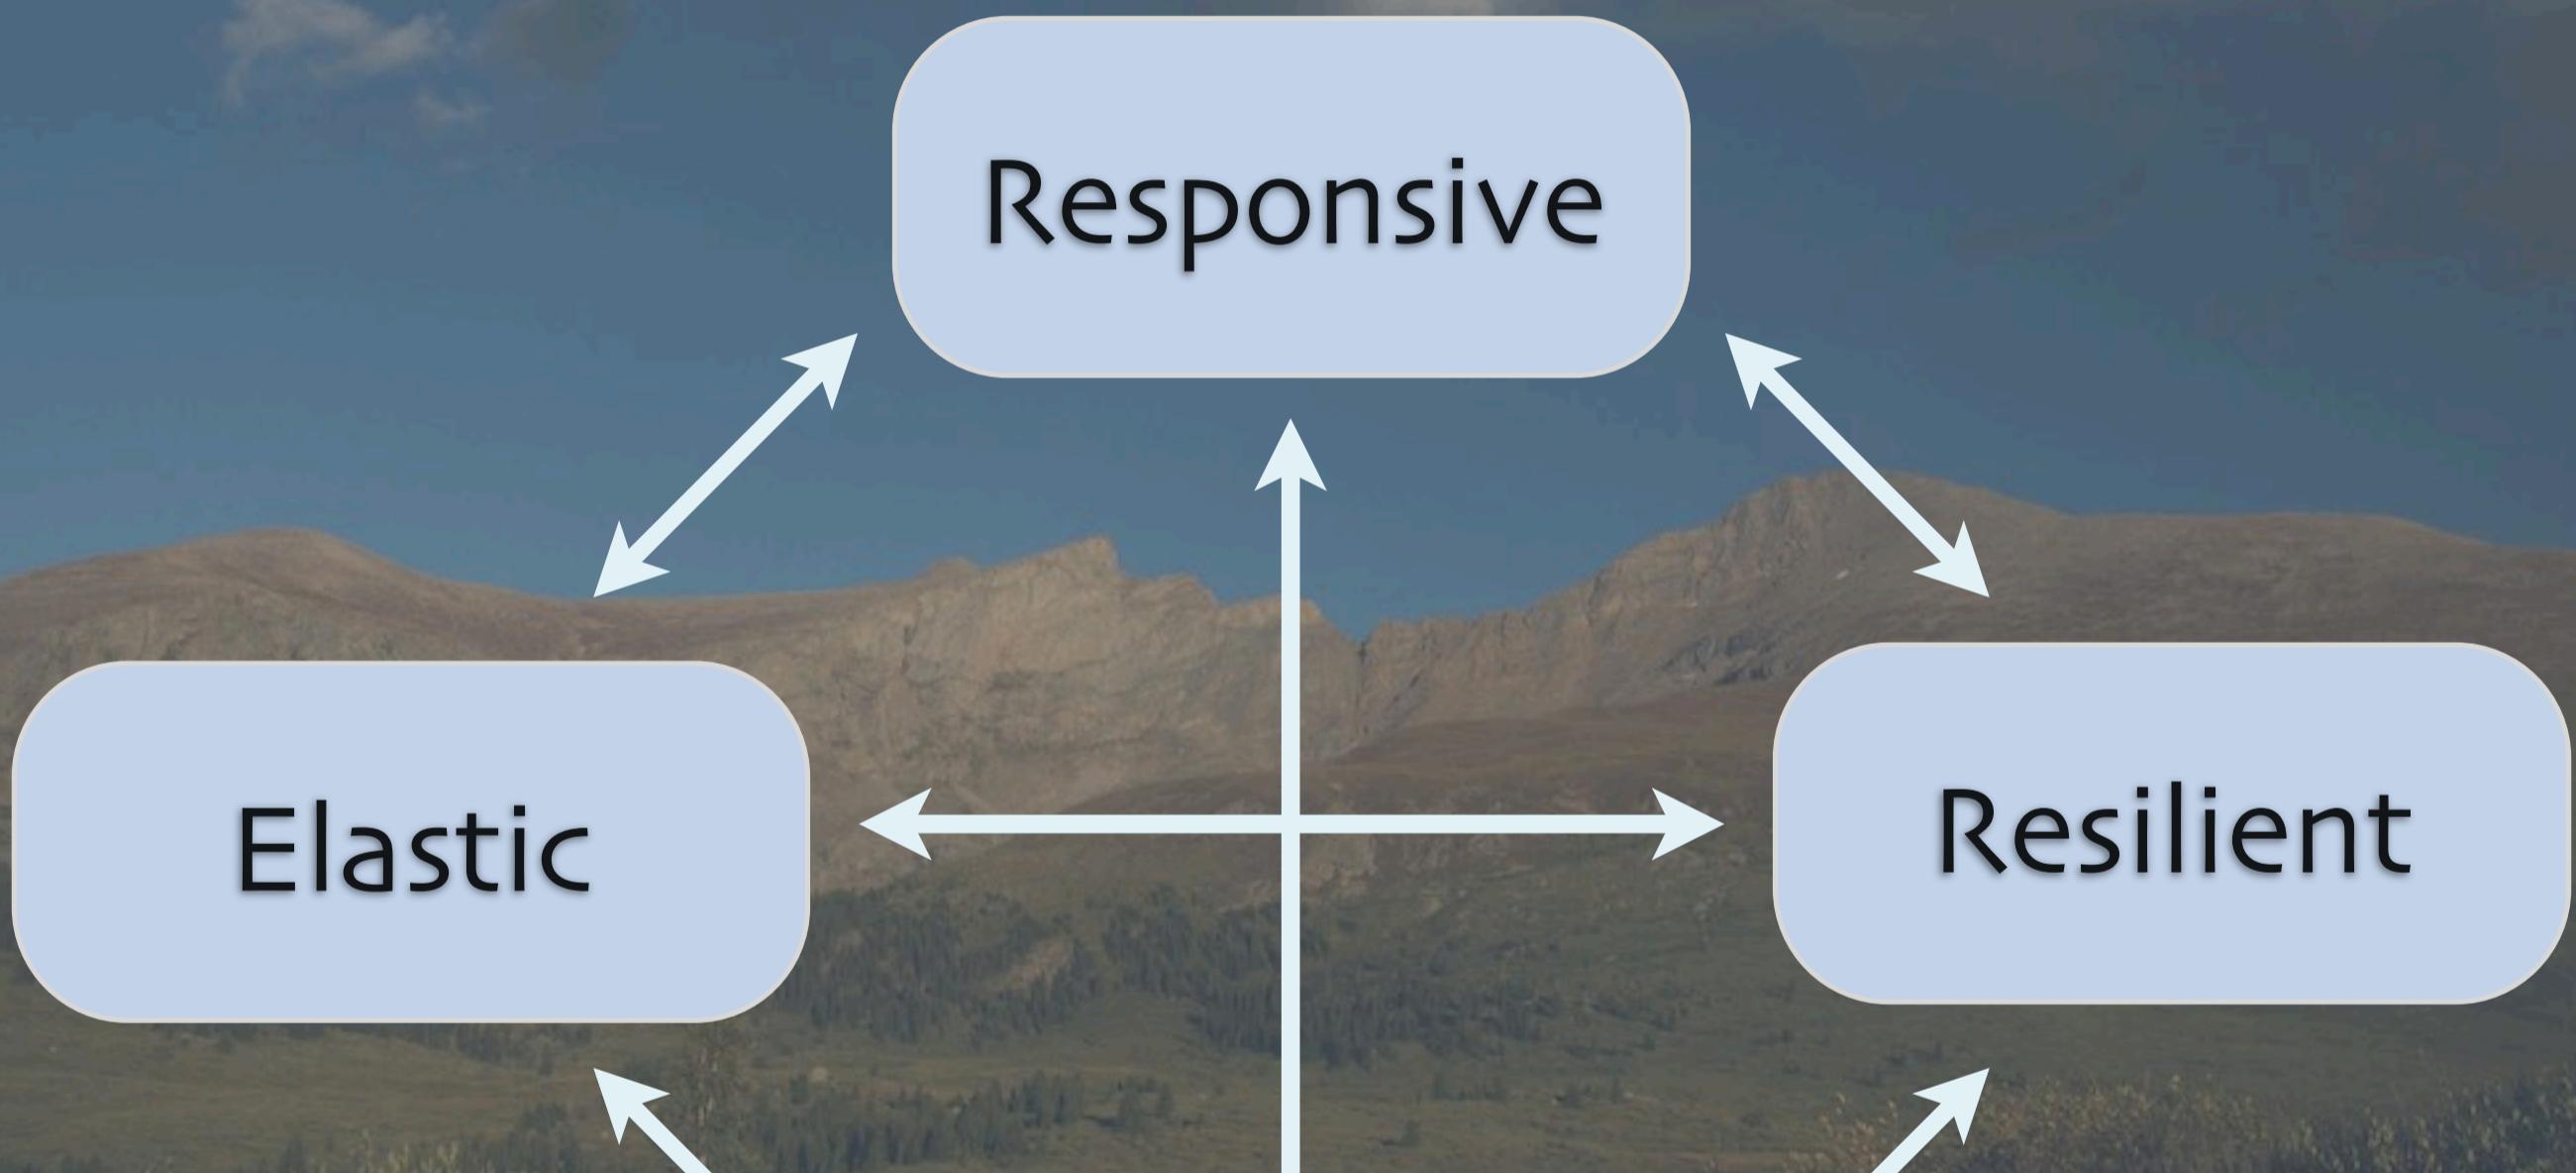
24

Tuesday, March 17, 15

While Typesafe's Jonas Bonér was one of the originators of the RM, other originators and contributions include experts in many companies and specialties.



Requests or commands
require timely responses.





Responsive



Responsive

Cornerstone of
usability and utility.



Responsive

Requires rapid
detection of errors
and quick responses.



Responsive

Requires predictable
response times
and quality of service.



Responsive

Awareness of time
is first class.

Example:

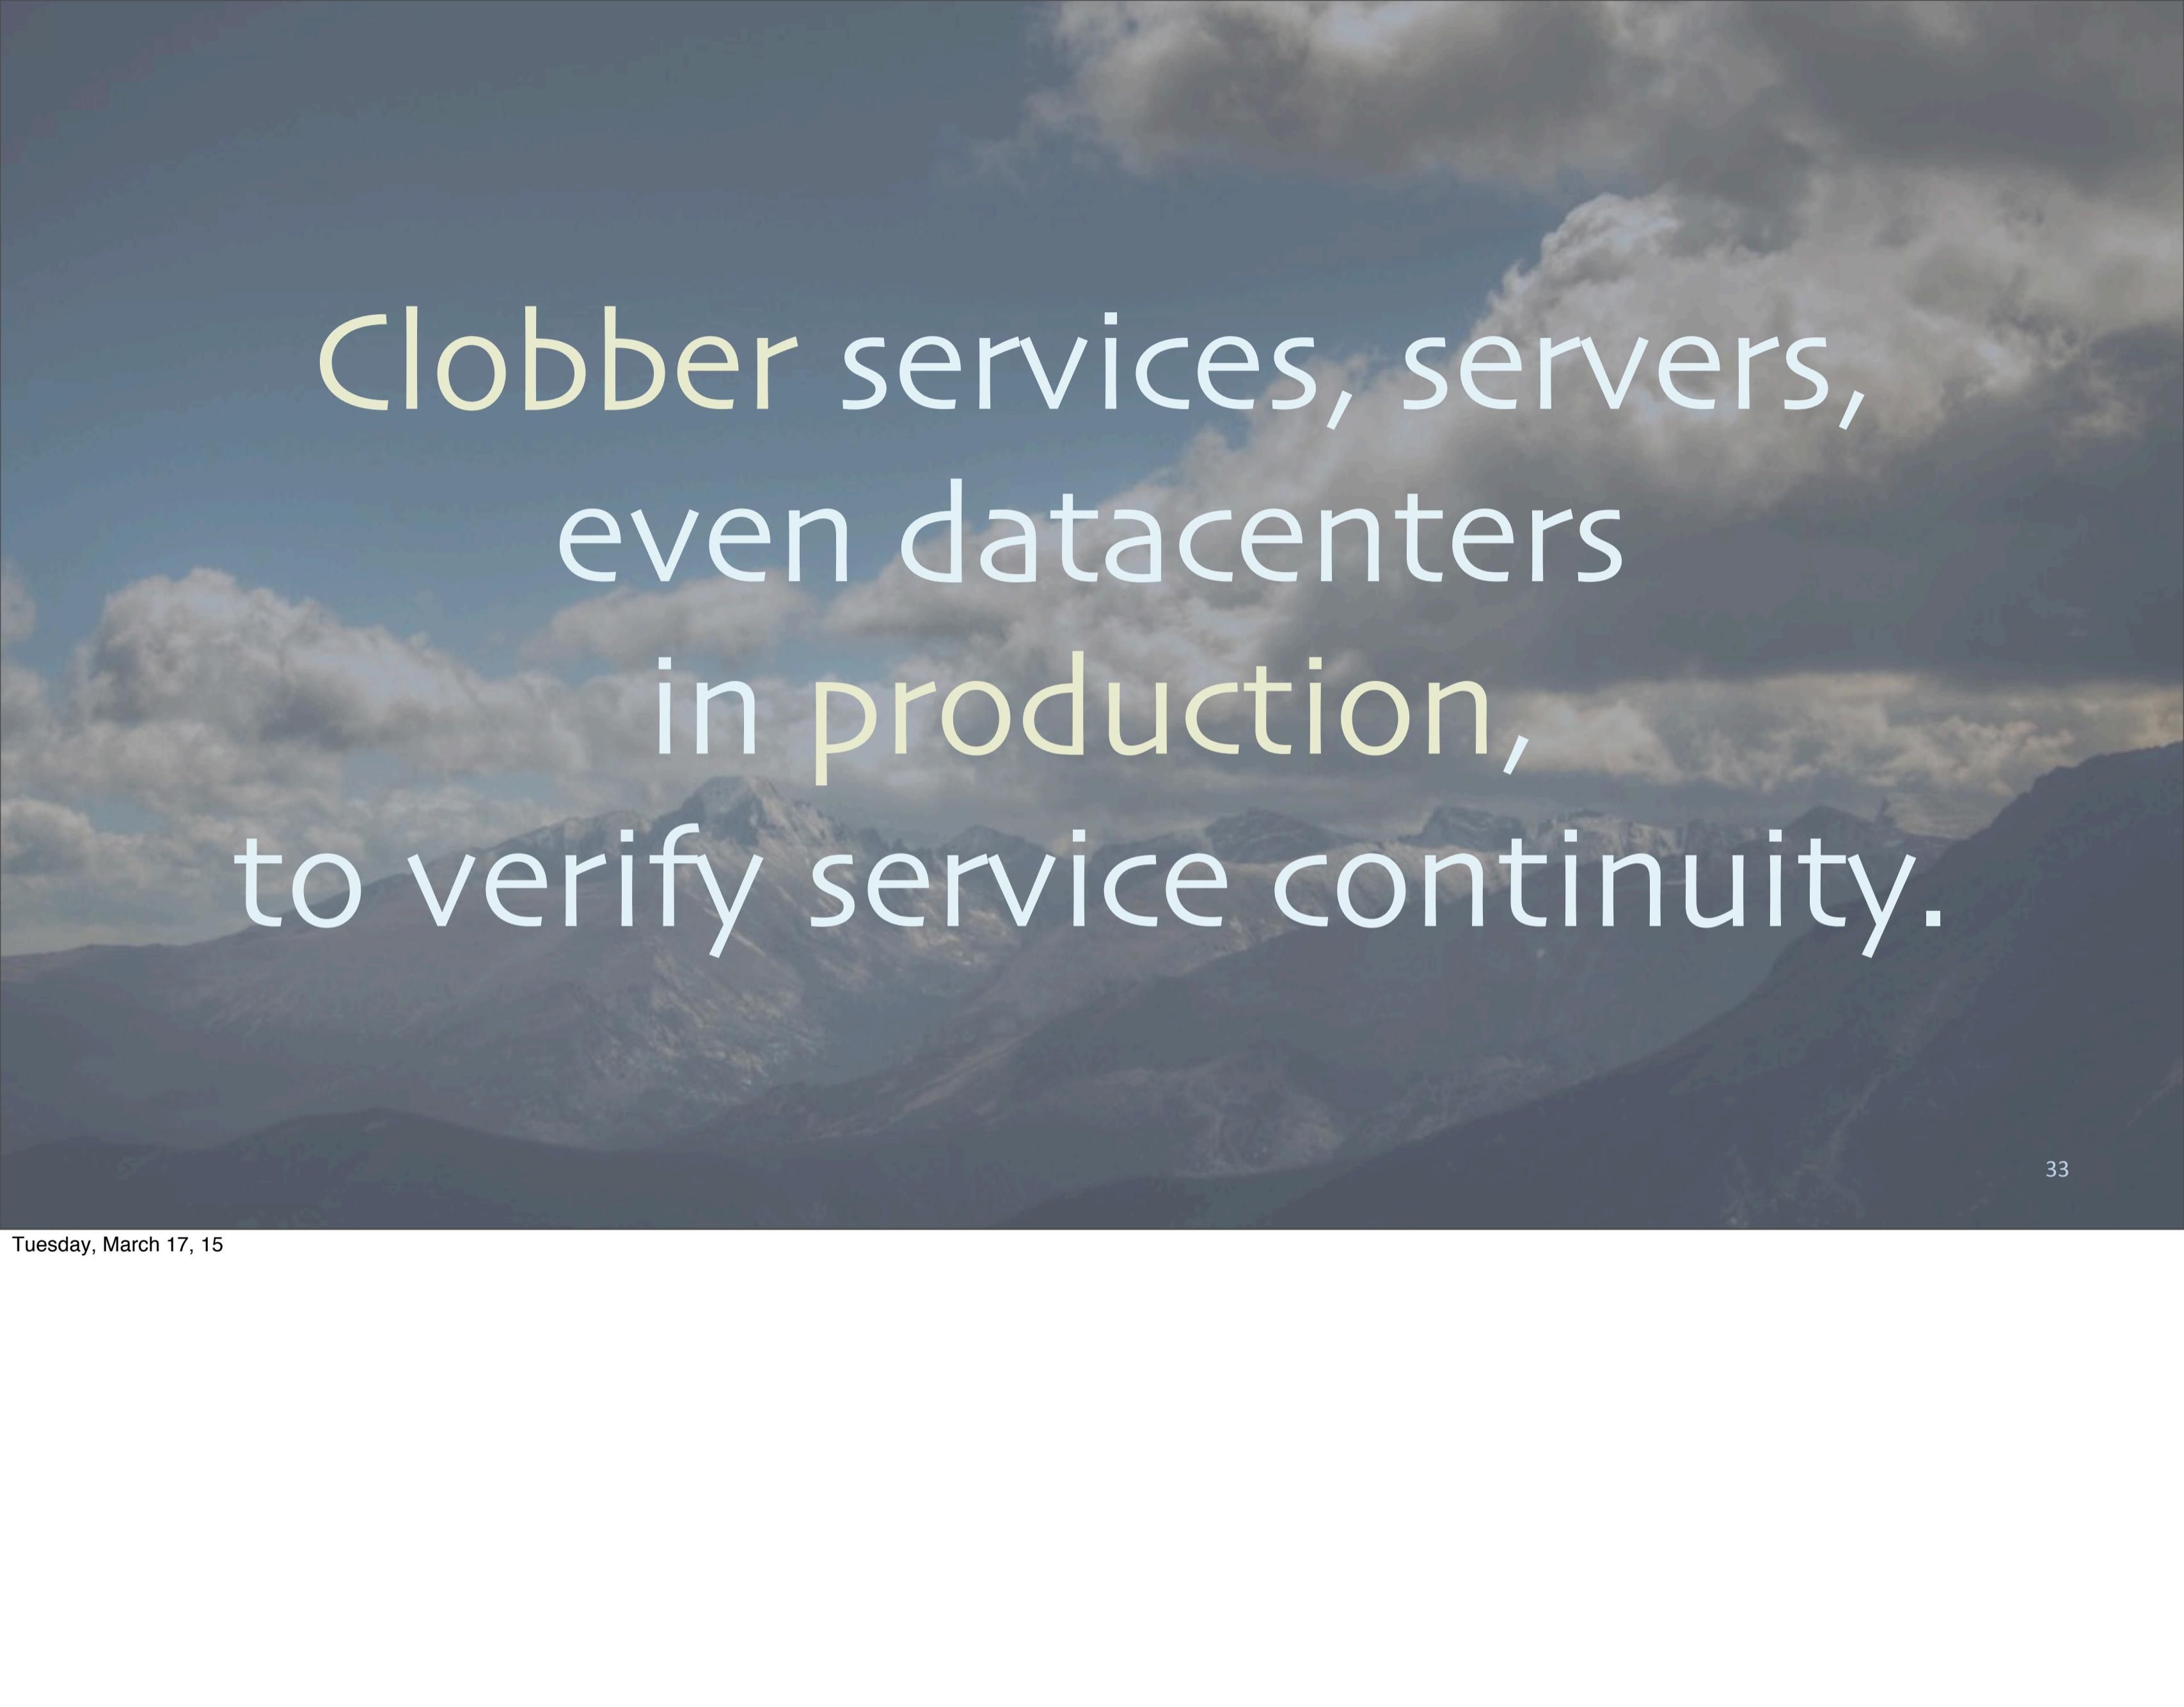
Netflix Simian Army



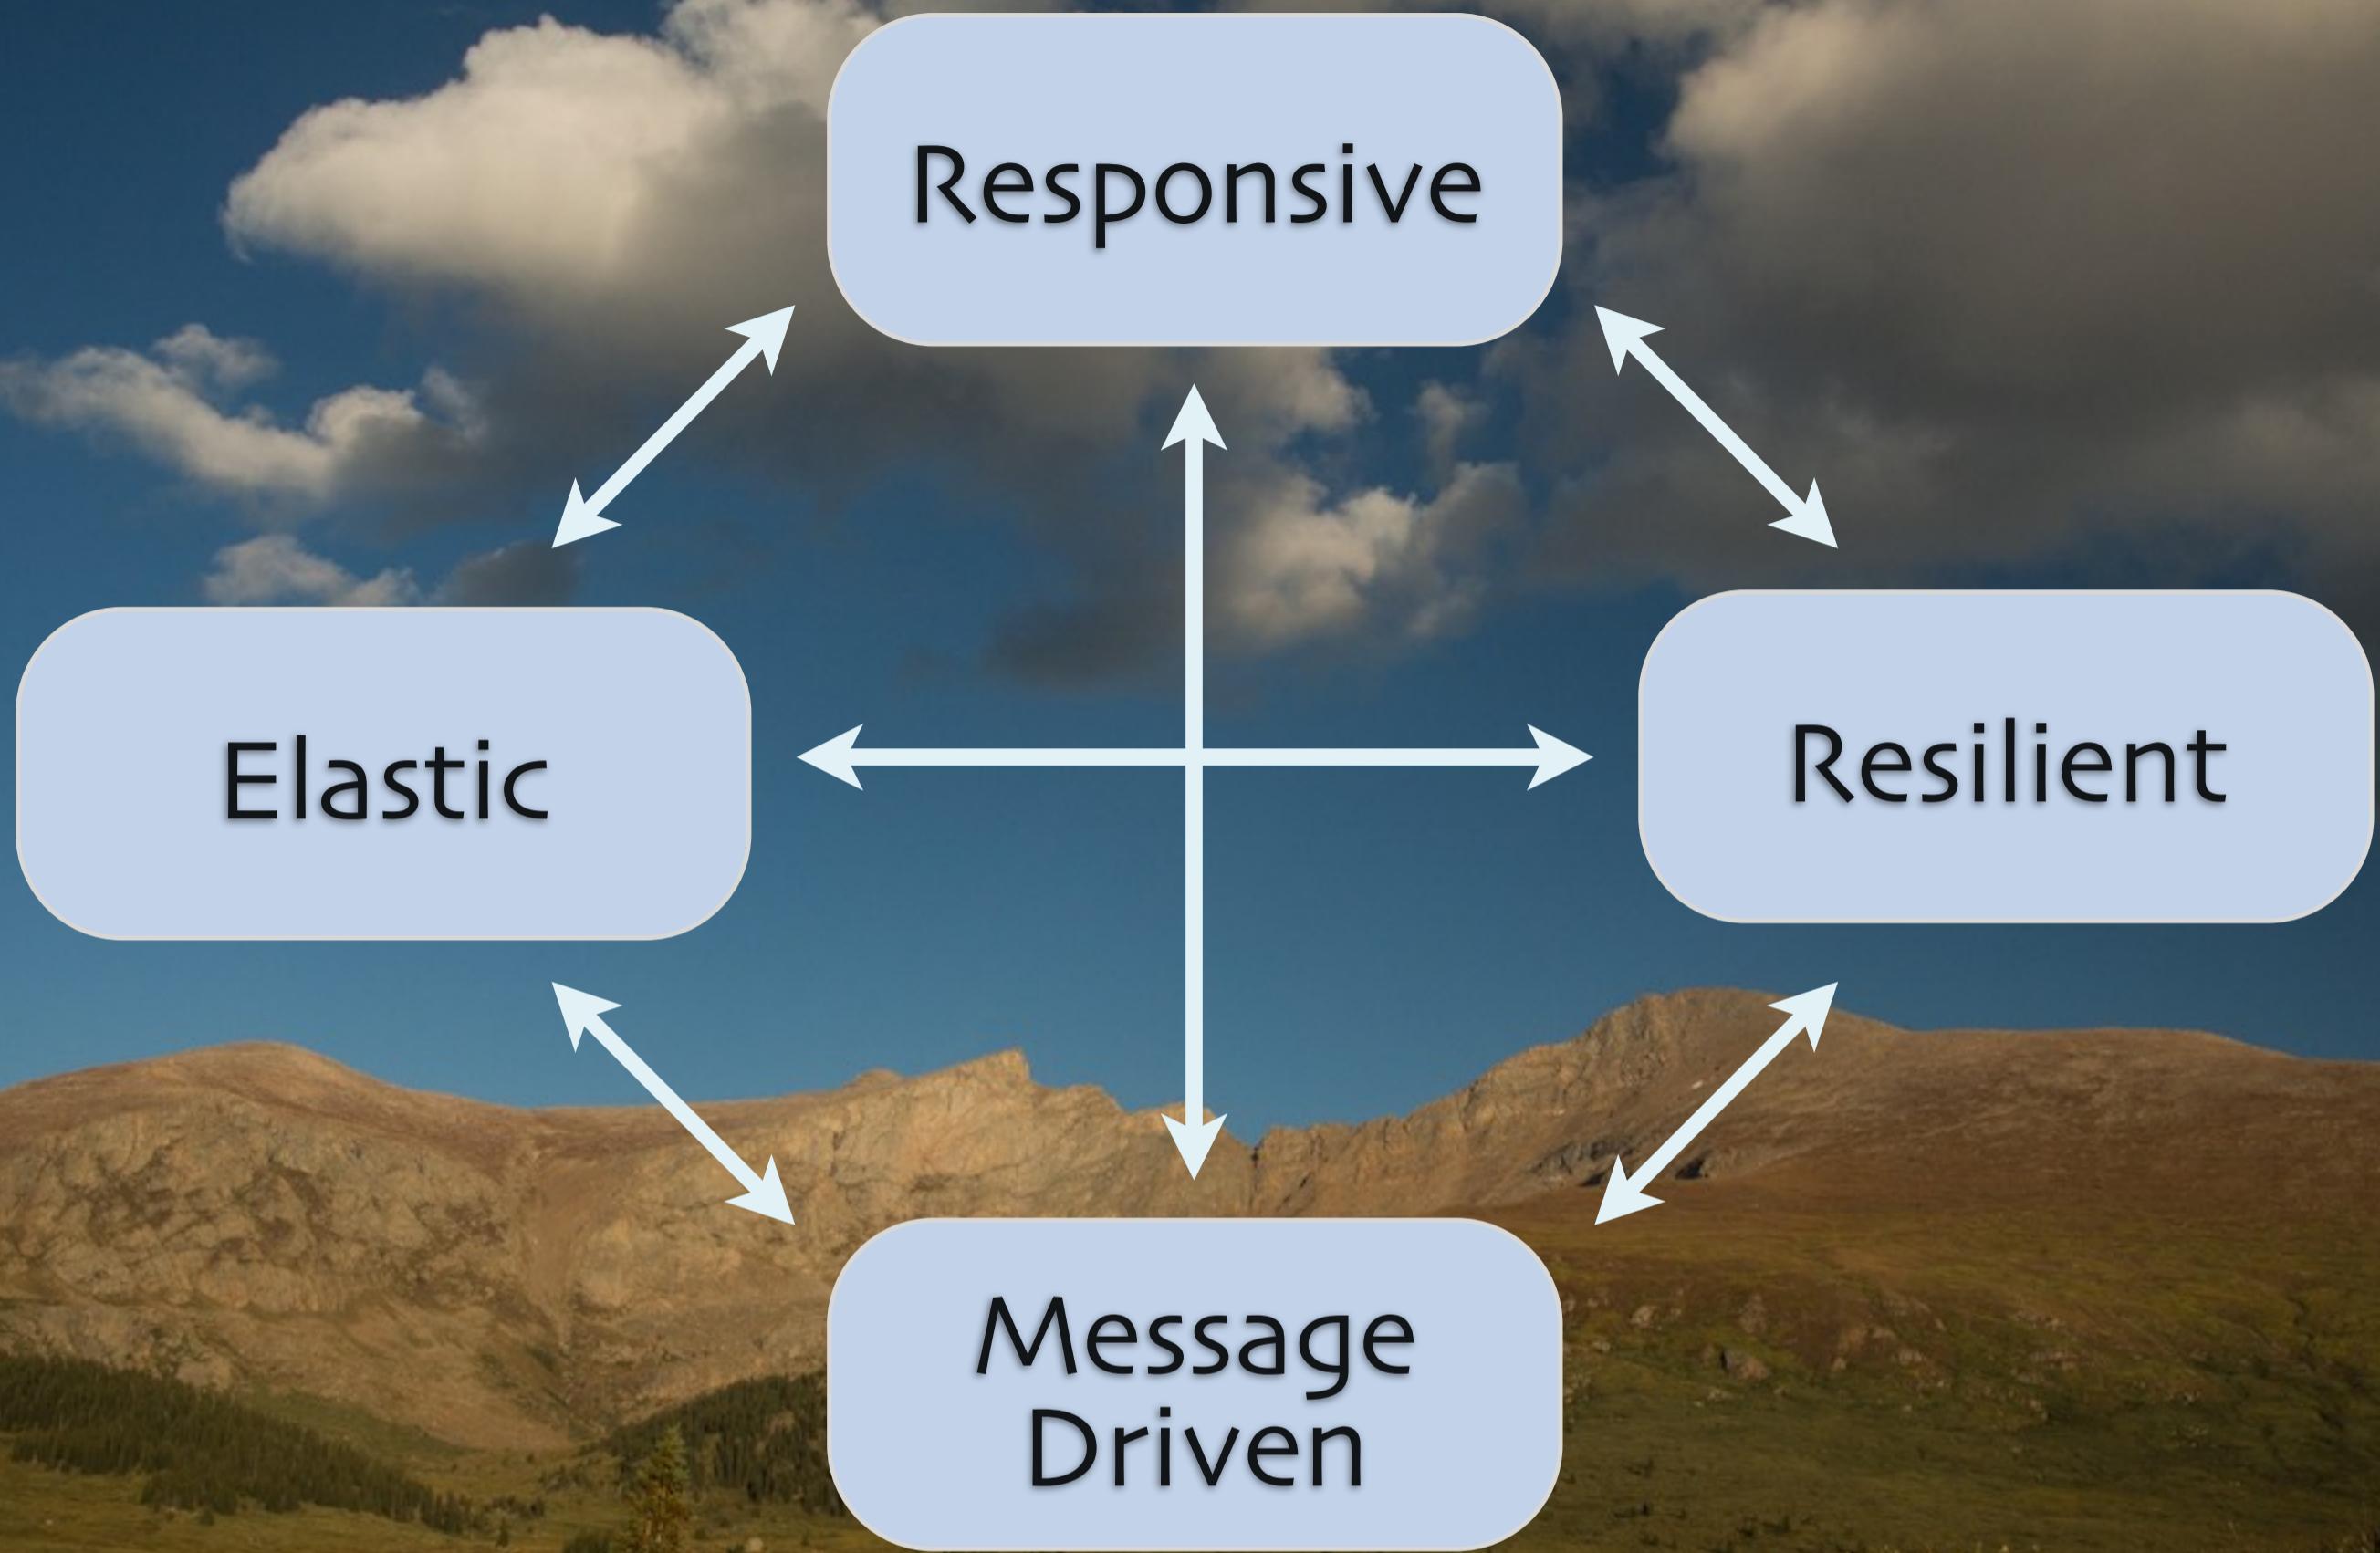
32

Tuesday, March 17, 15

Image from <http://devops.com/features/netflix-the-simian-army-and-the-culture-of-freedom-and-responsibility/>



Clobber services, servers,
even datacenters
in production,
to verify service continuity.





35

Tuesday, March 17, 15

Truly resilient systems must make failures first class citizens, in some sense of the word, because they are inevitable when the systems are big enough and run long enough.



Resilient



Resilient

Failure is
not disruptive.



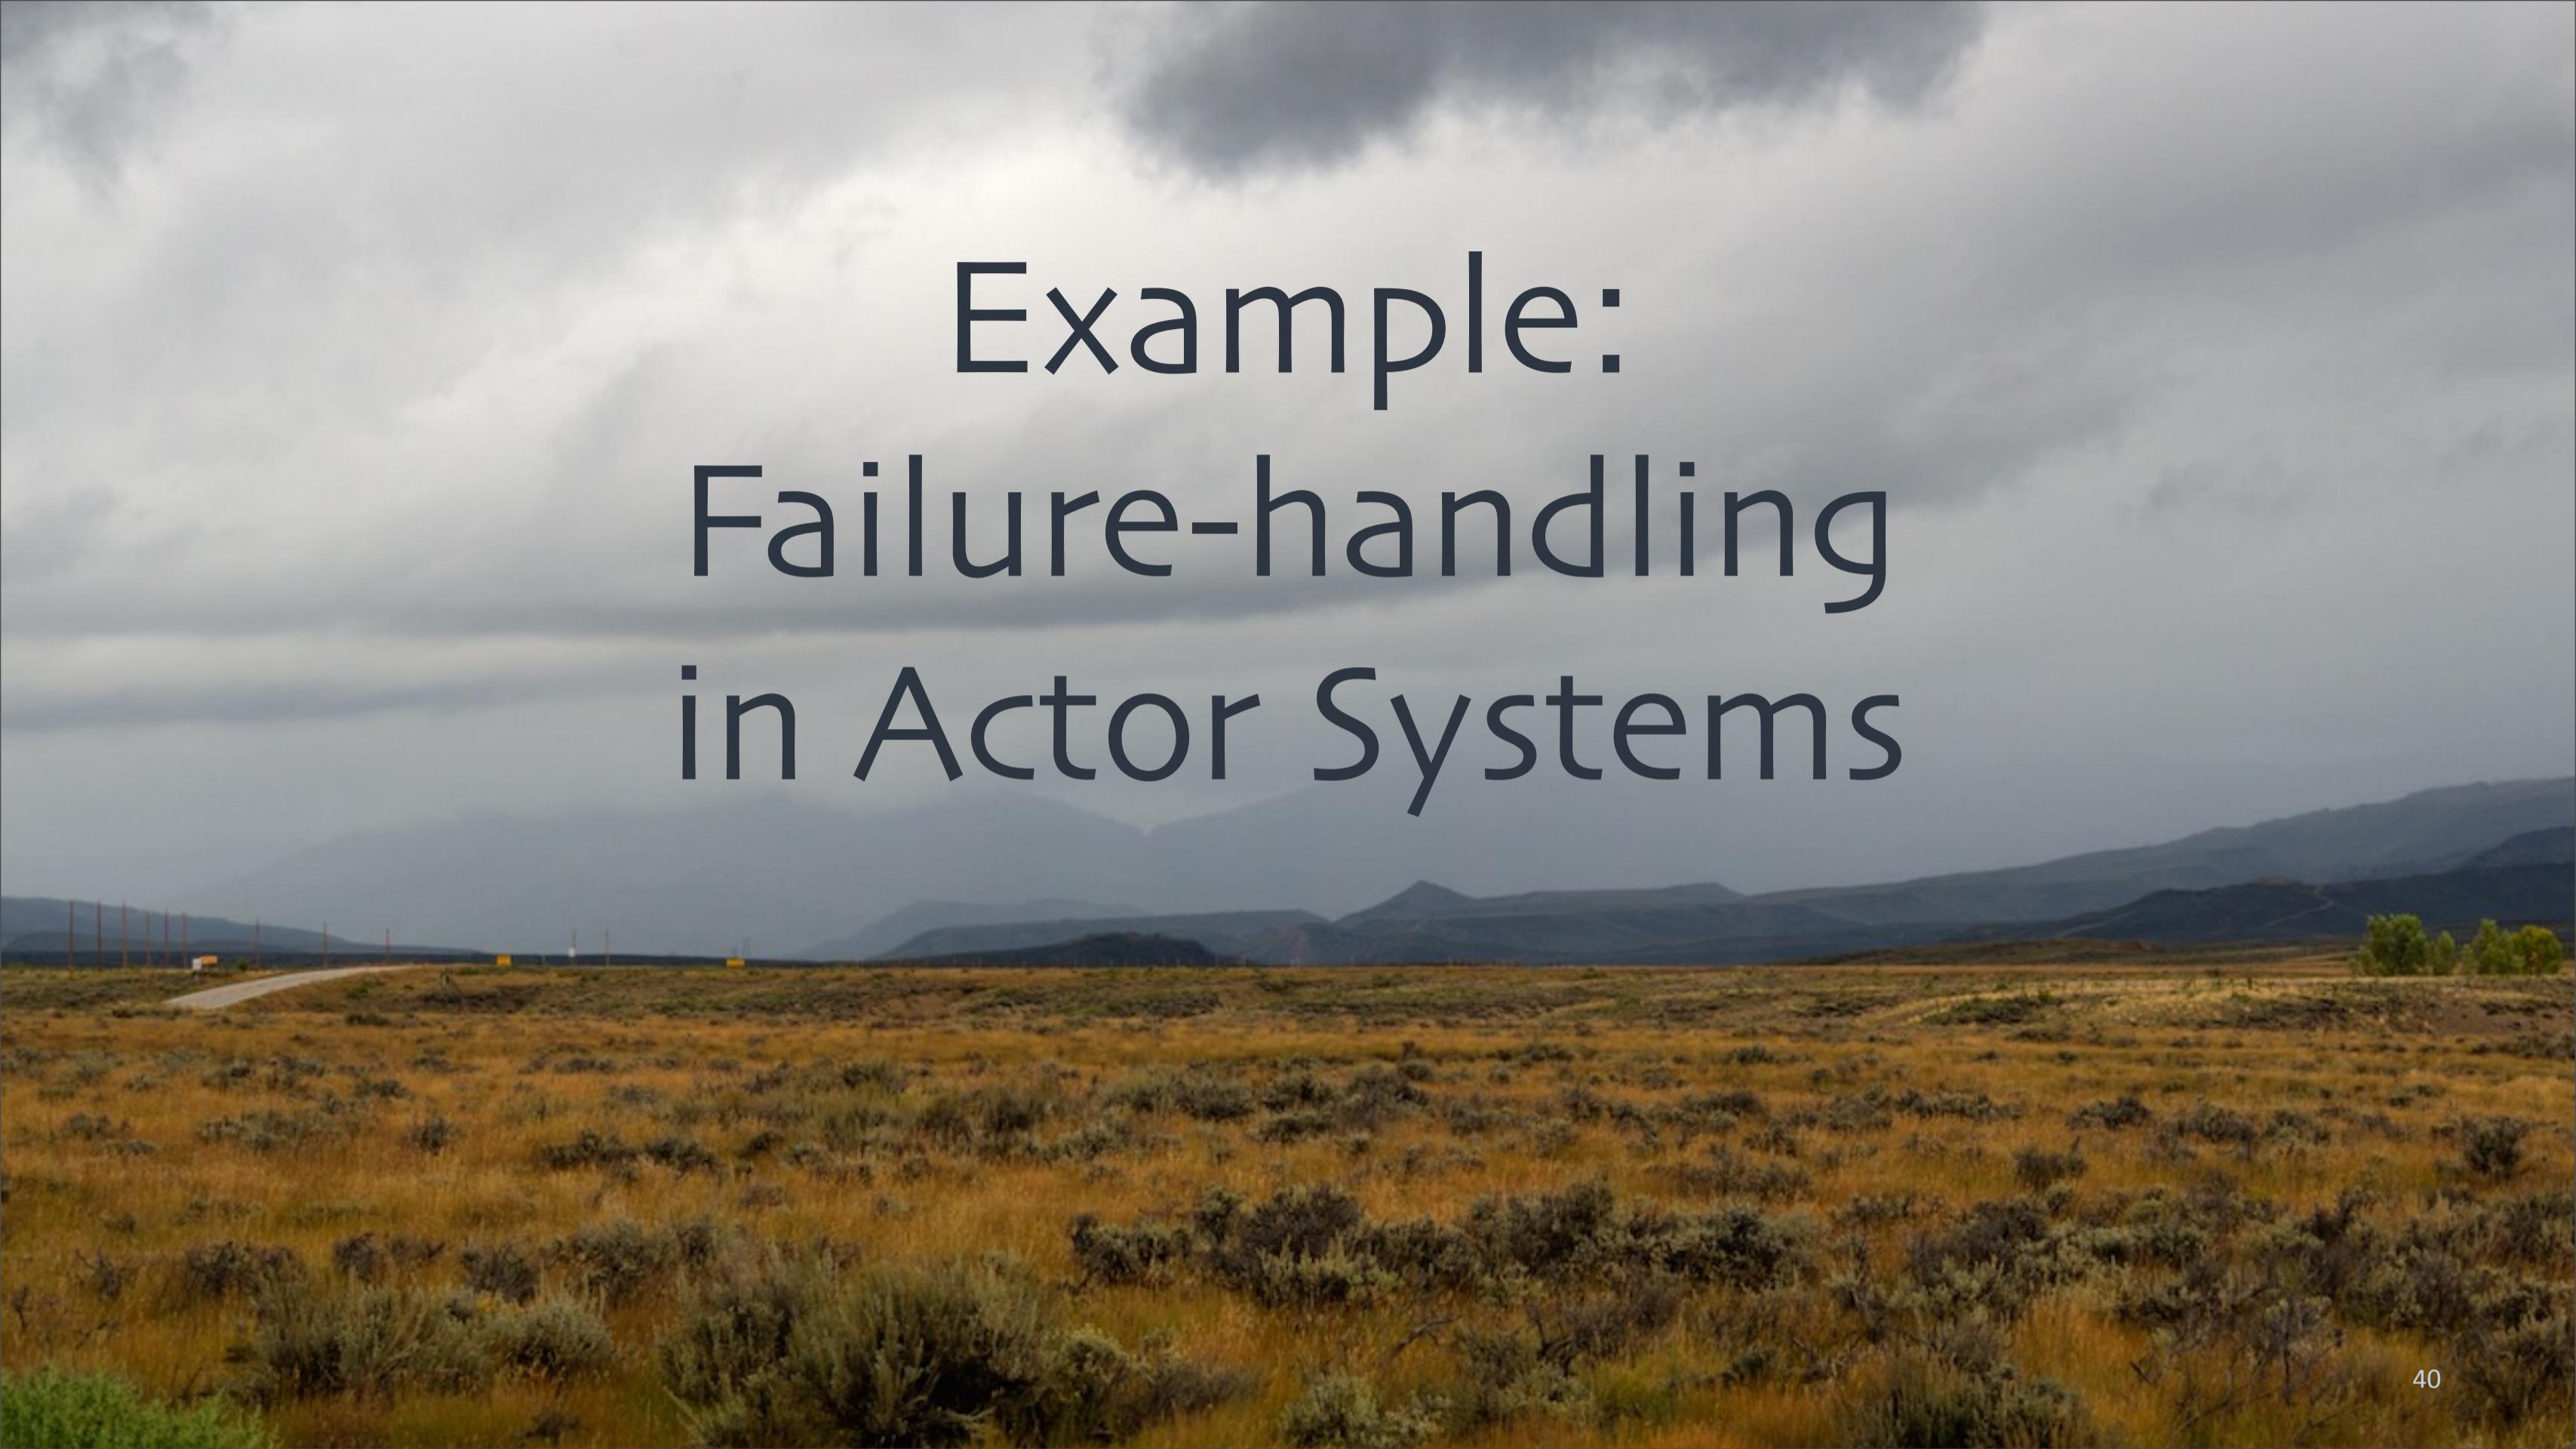
Resilient

Failure is
first class.

A scenic landscape featuring a range of mountains in the background, covered with green forests. In the foreground, there's a field of tall grass. The sky is filled with large, white, billowing clouds.

Resilient

Requires replication,
containment, isolation,
and delegation.



Example: Failure-handling in Actor Systems

40

Tuesday, March 17, 15

The Netflix Simian Army could also be cited here.
We'll come back and fill in details of Actor systems shortly. For now, let's focus on error handling.



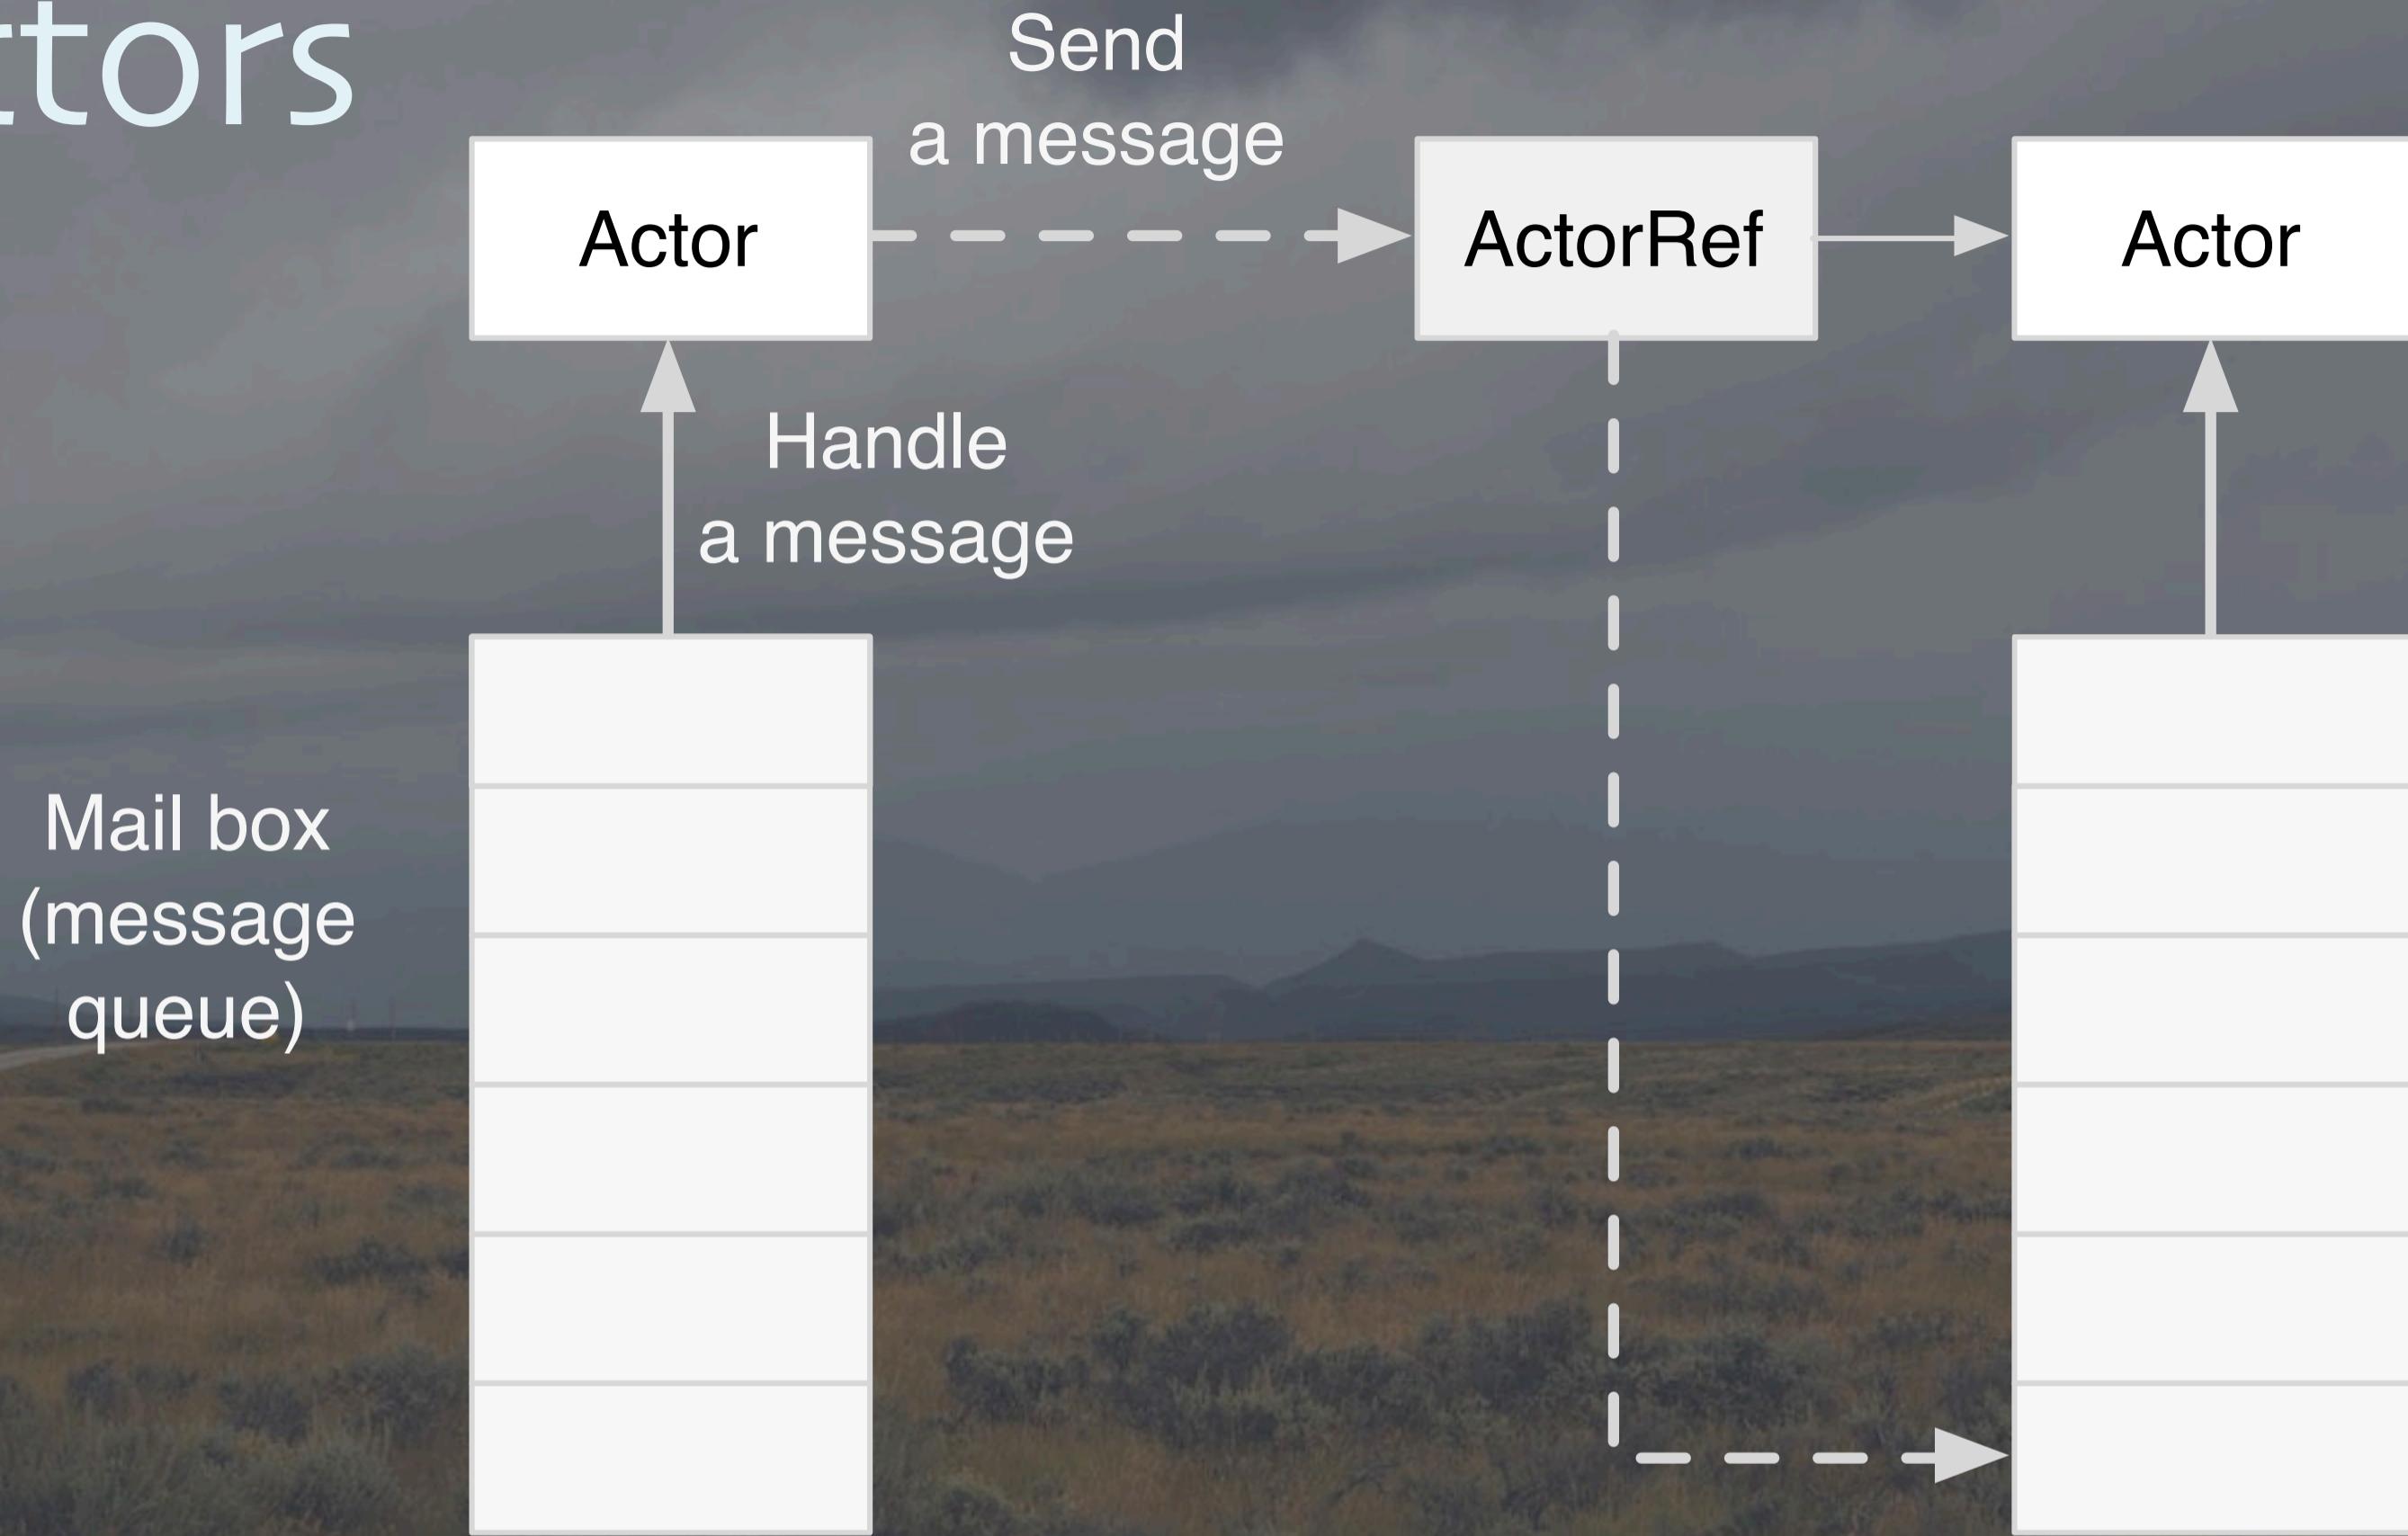
Let it Crash!

41

Tuesday, March 17, 15

Rather than attempt to recover from errors inside the domain logic (e.g., elaborate exception handling), allow services to fail, but with failure detection and reconstruction of those services, plus failover to other replicas.

Actors

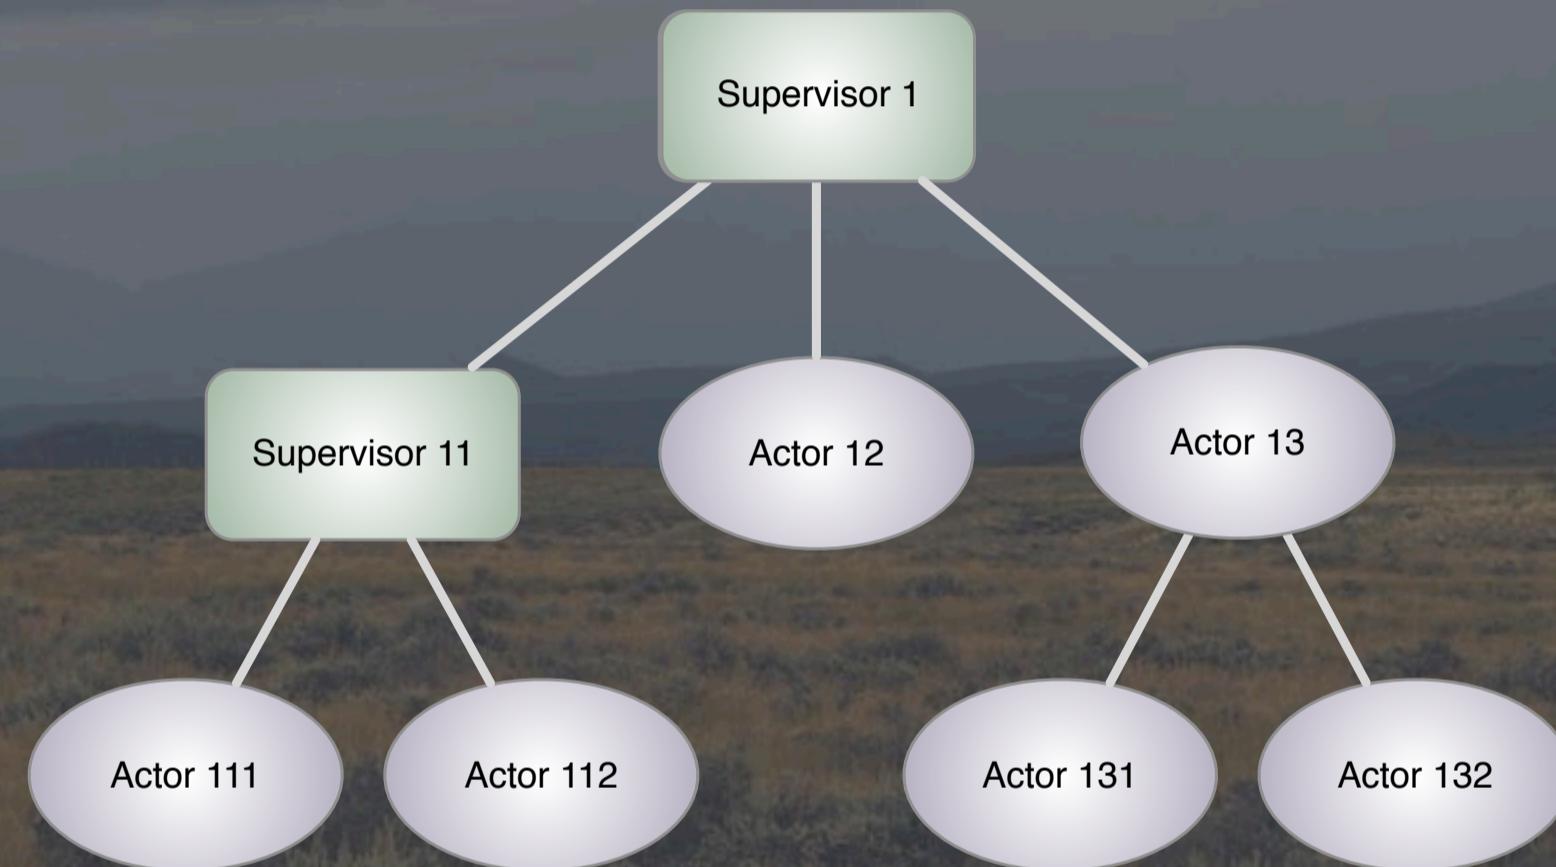


42

Tuesday, March 17, 15

Actors are similar to objects in Smalltalk and similar systems; autonomous agents with defined boundaries that communicate through message passing. Actors, though process each message in a threadsafe way, so they are great for concurrency. (This diagram illustrates the Akka implementation – <http://akka.io>)

Erlang introduced supervisors. A hierarchy of actors that manage each “worker” actor’s lifecycle.

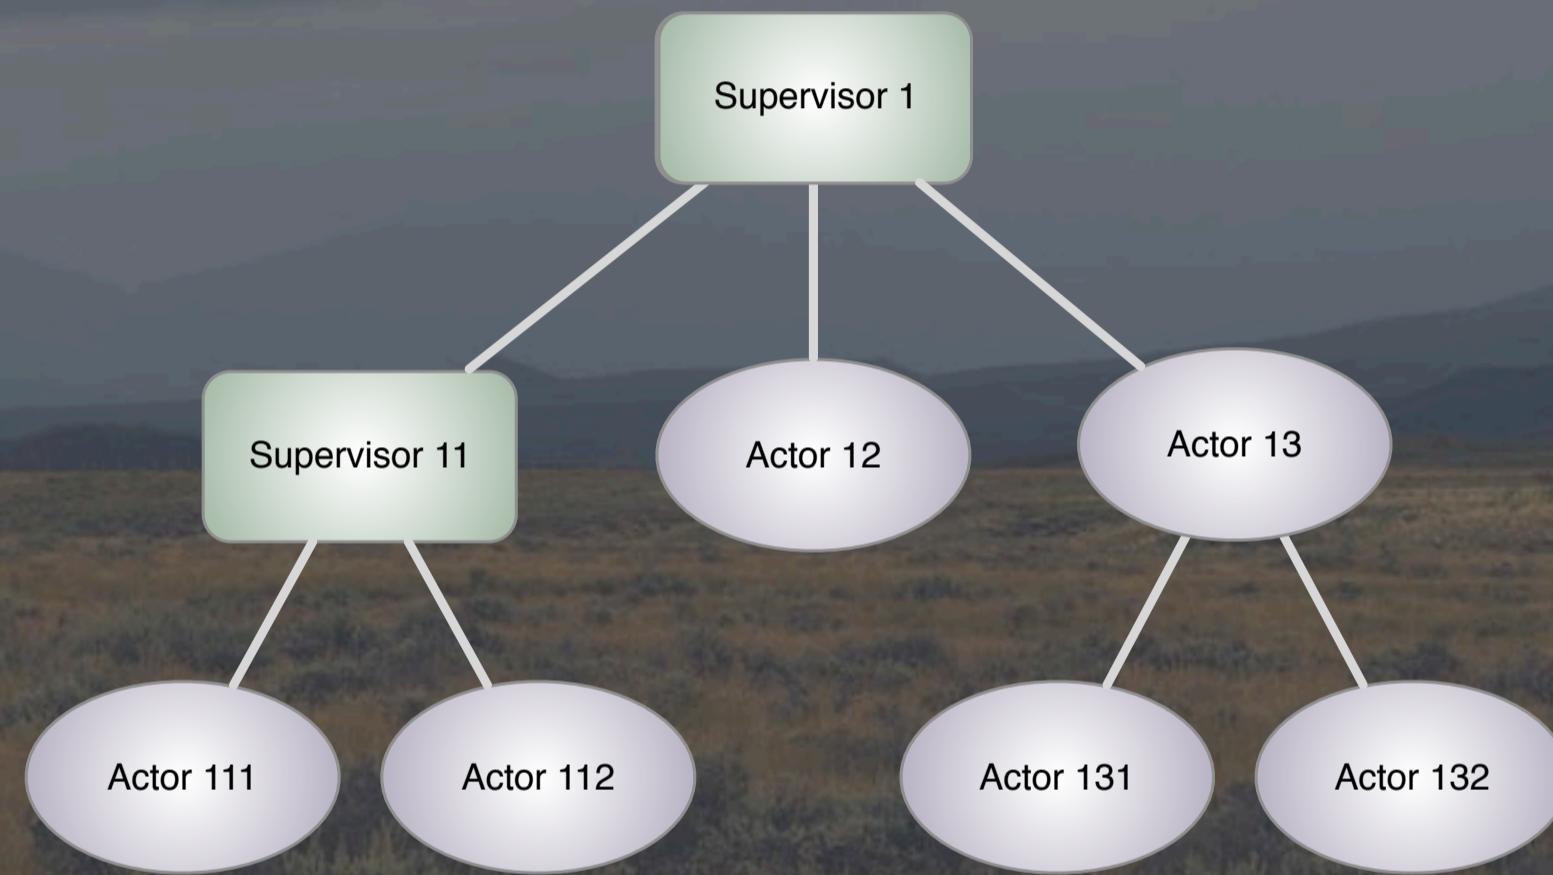


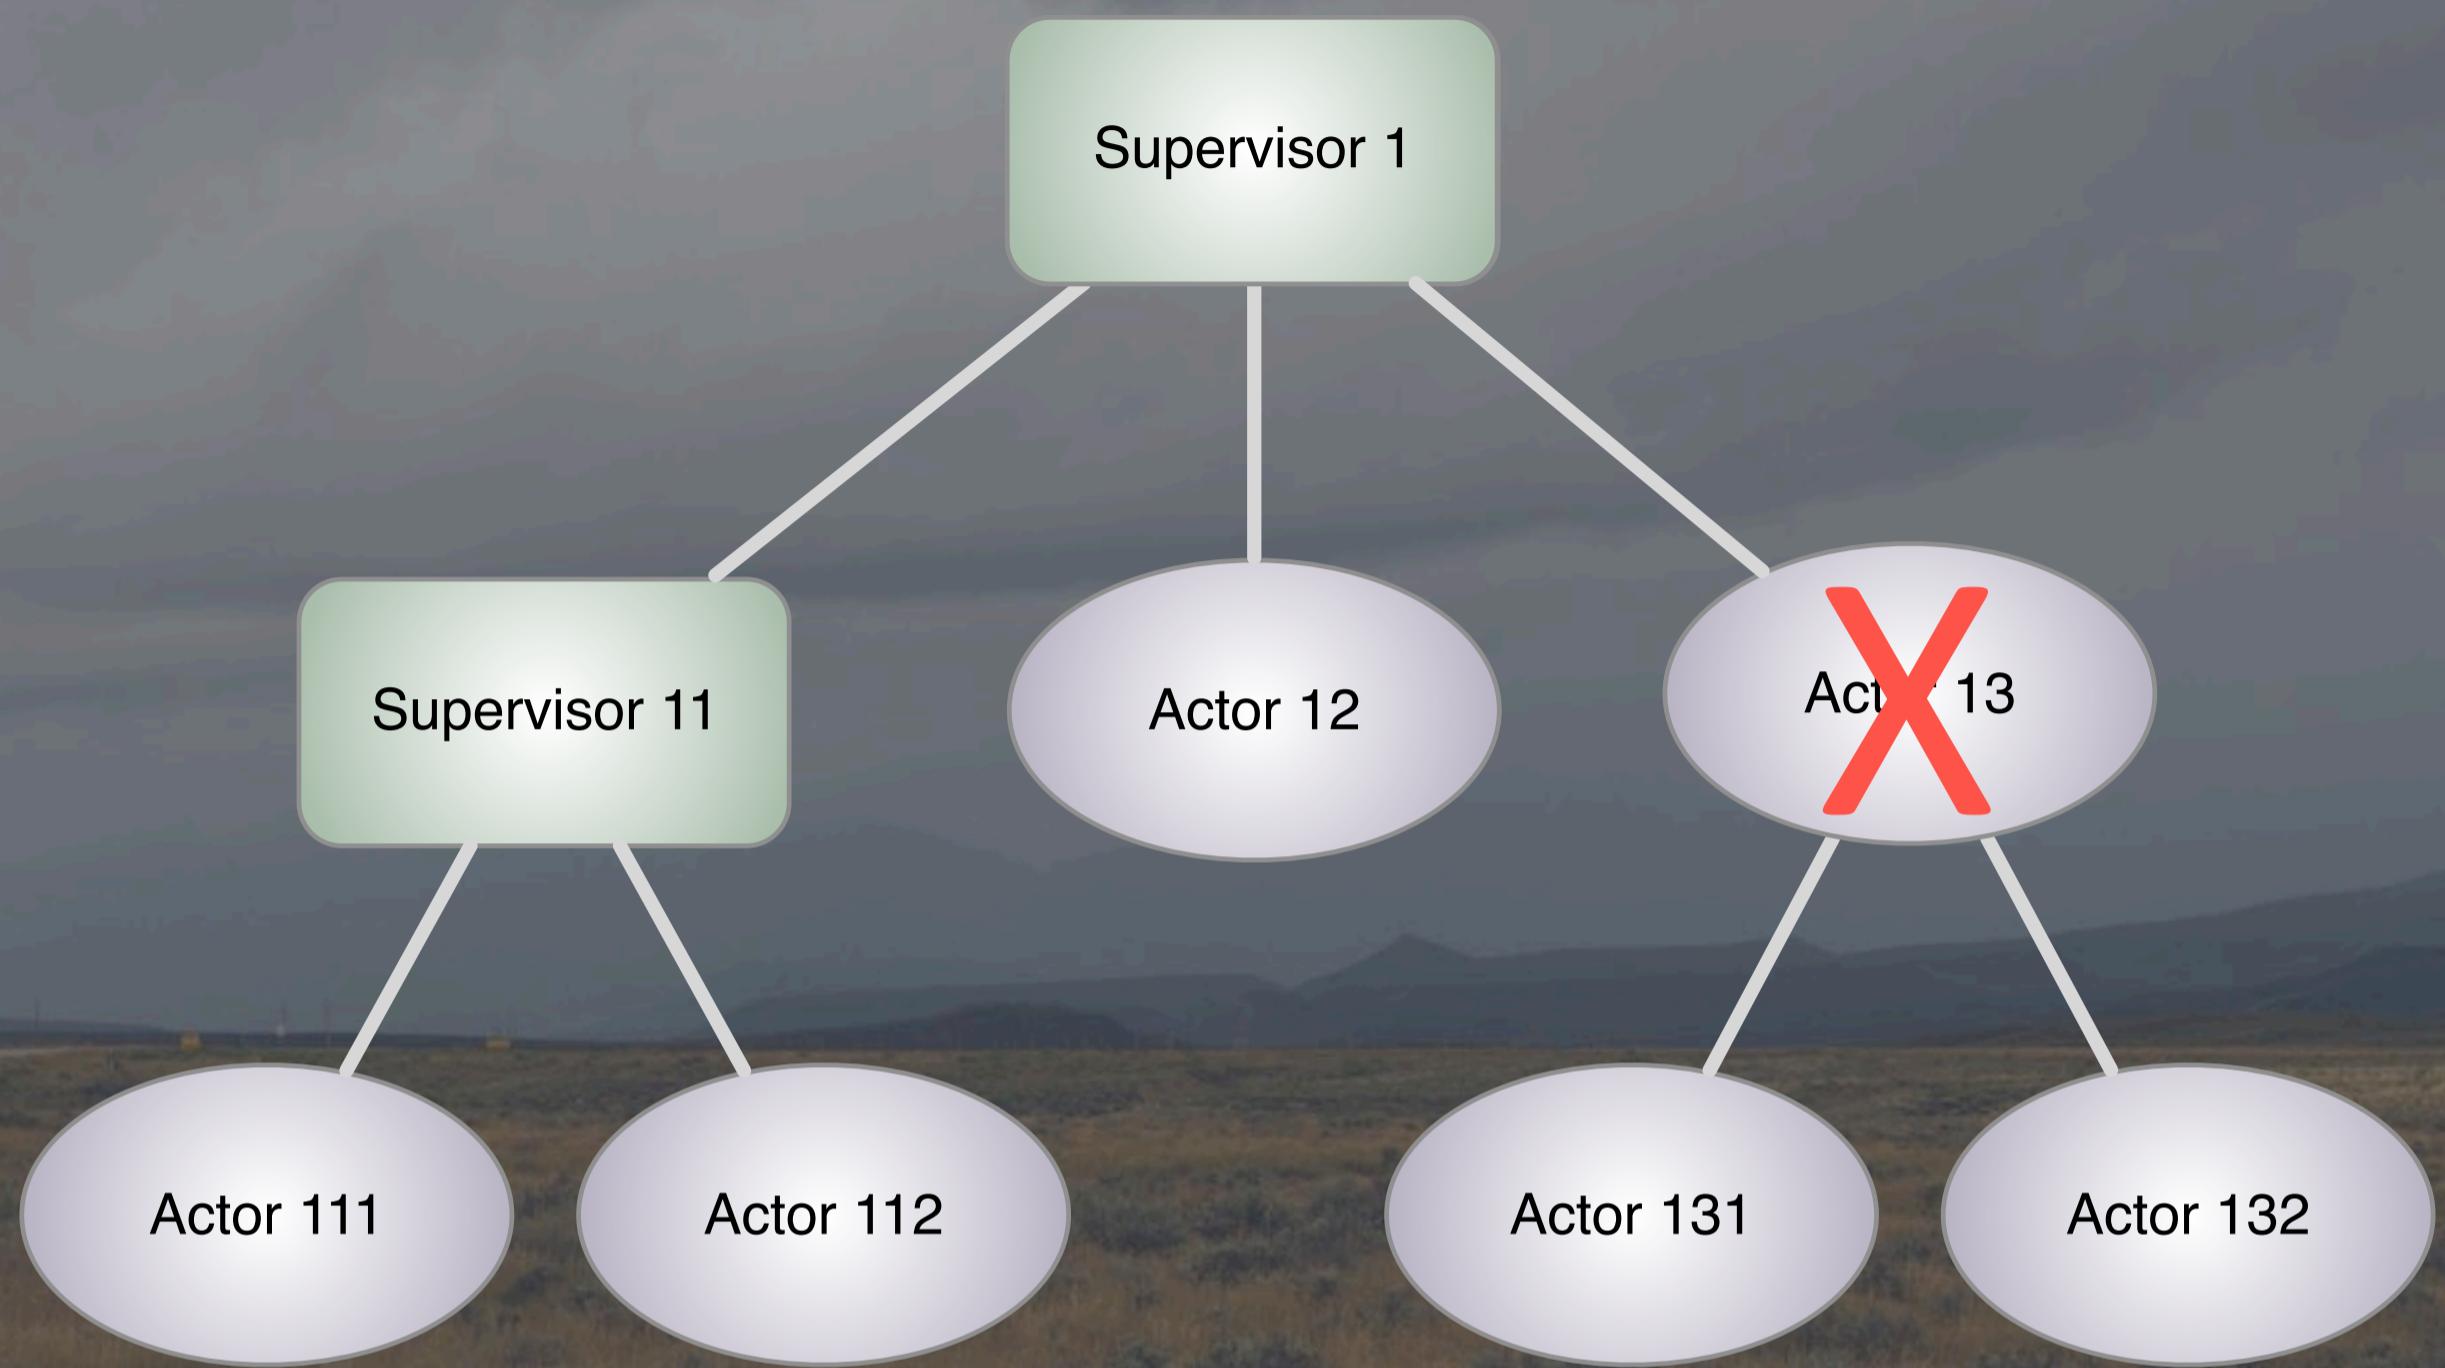
43

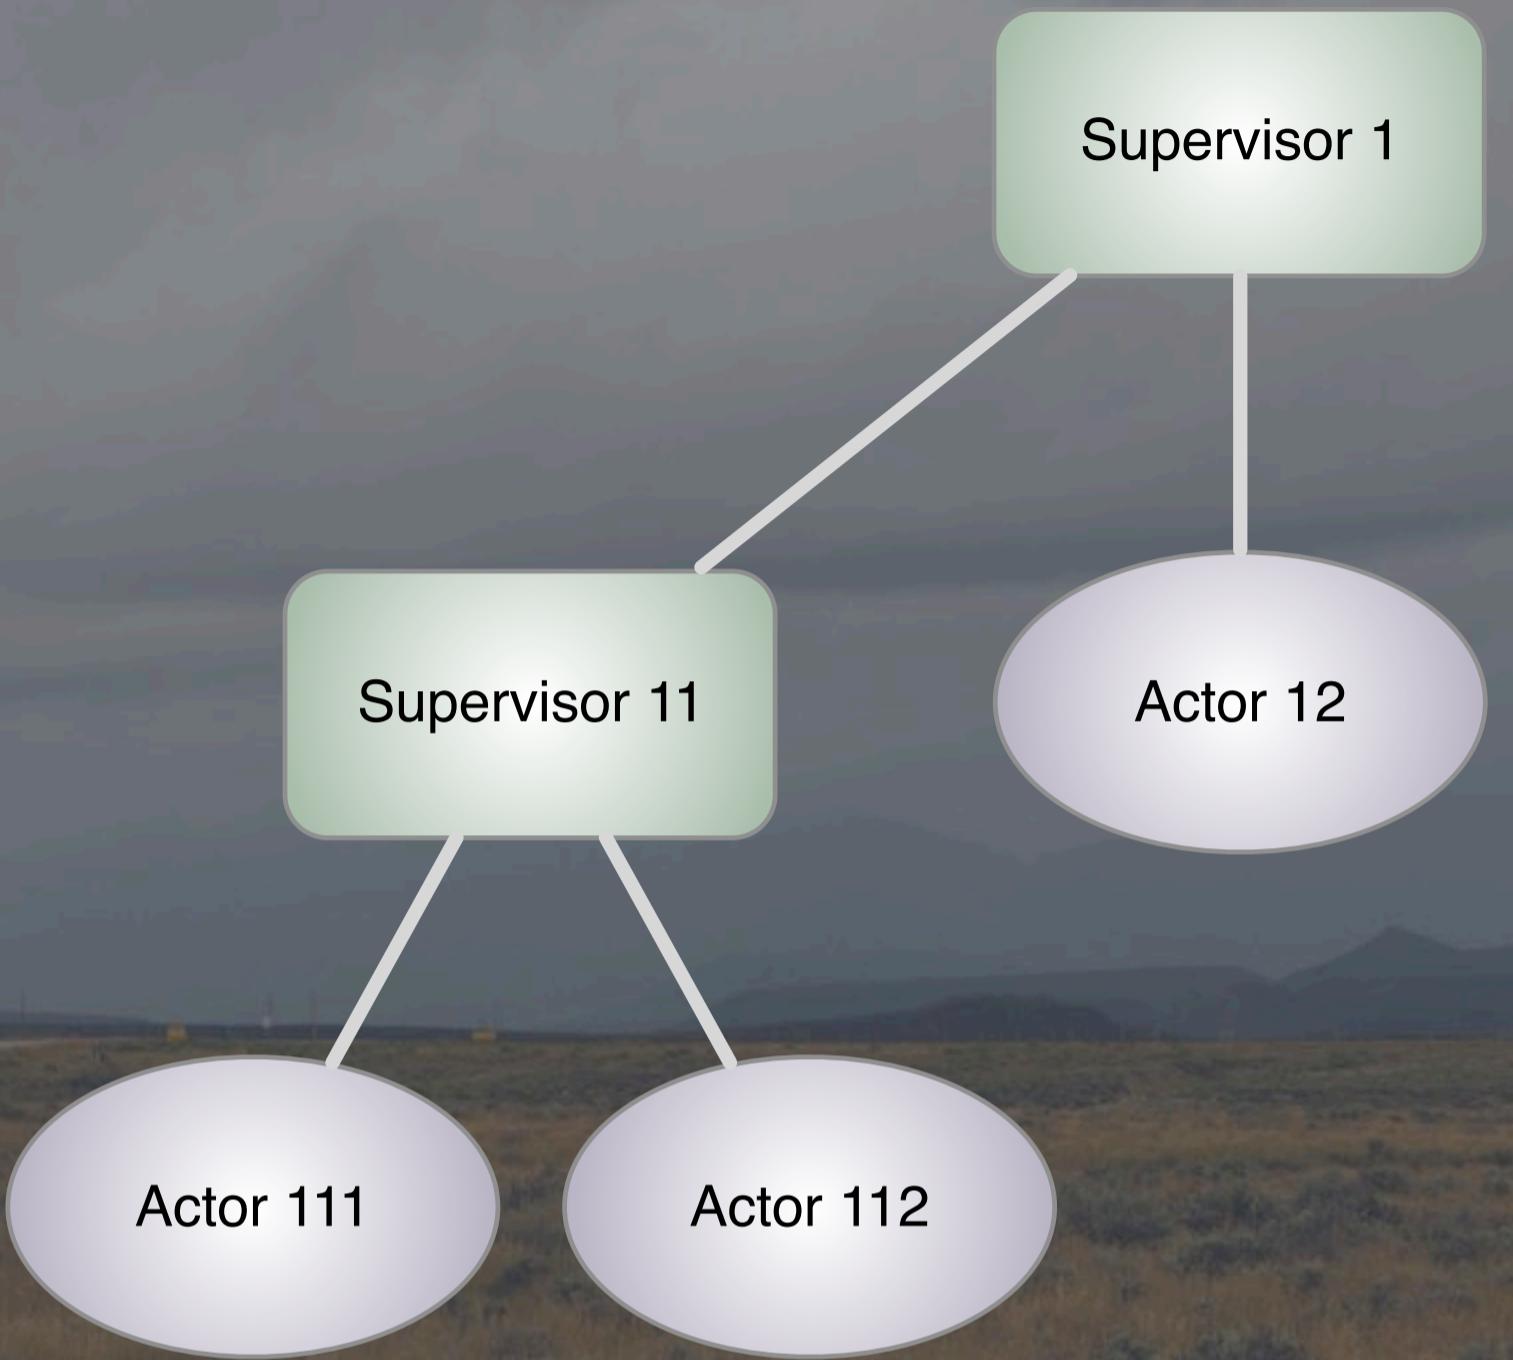
Tuesday, March 17, 15

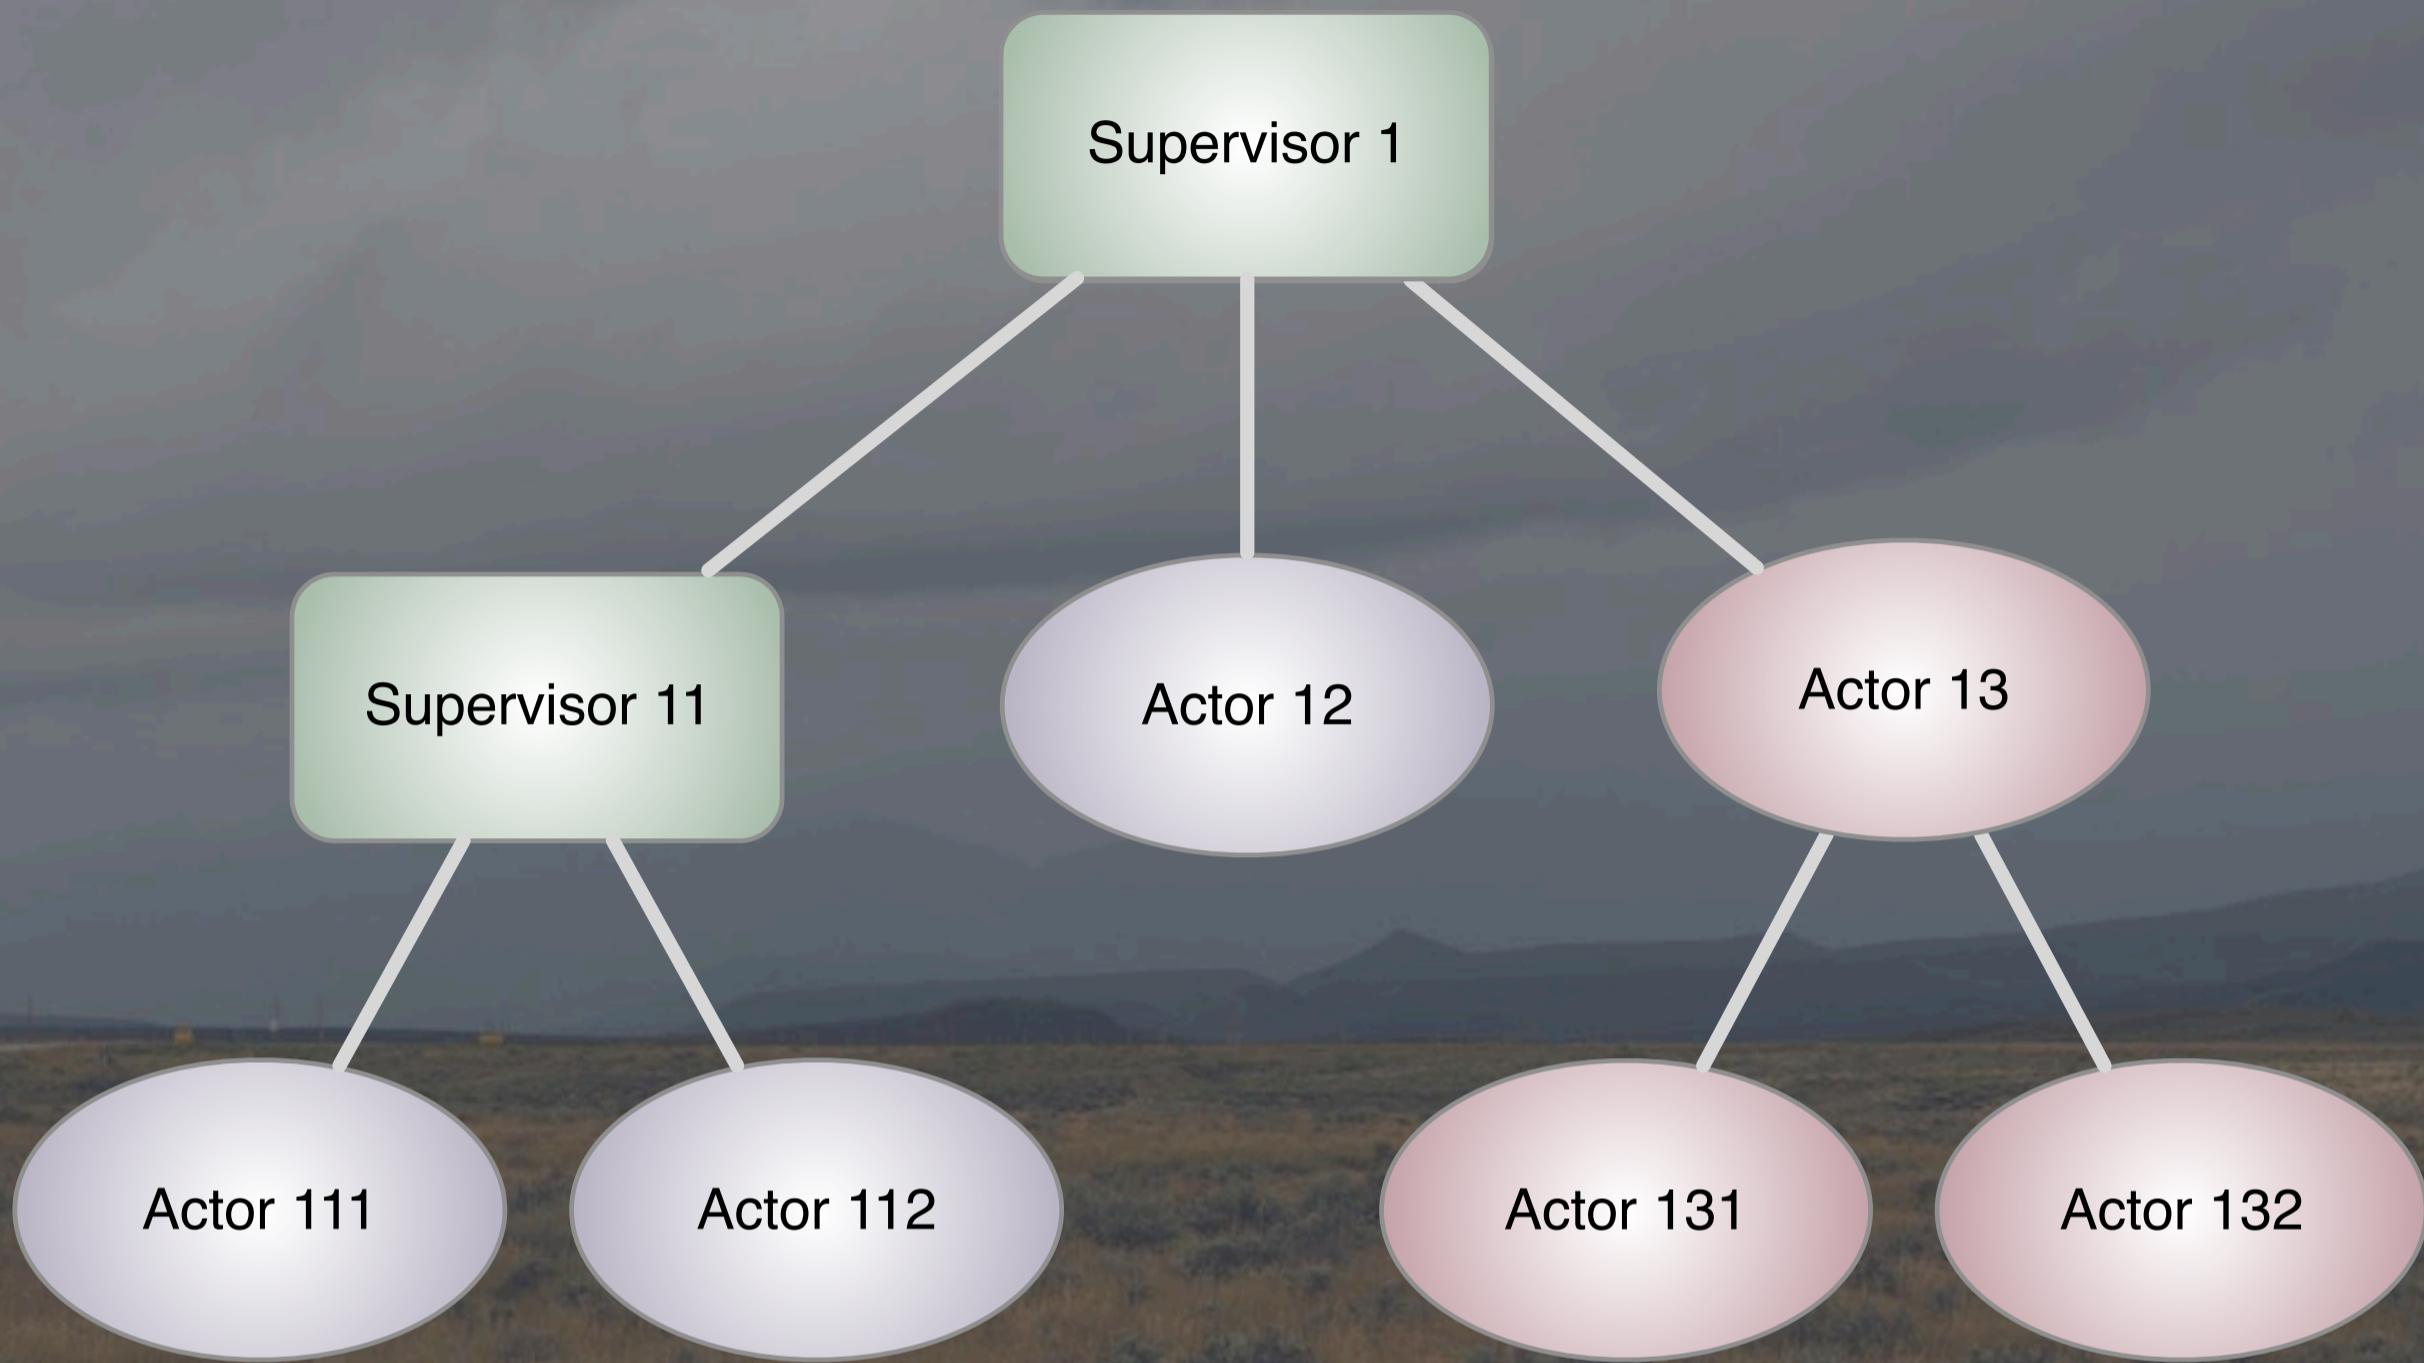
There are lots of so-called Actor systems, but be wary of them unless they have this sophisticated supervision model or something like it (even though the original Actor model of Hewitt, et al., didn’t include supervision like this...).

Generalizes nicely to distributed actor systems.











The most
sophisticated error recovery
in reactive systems.

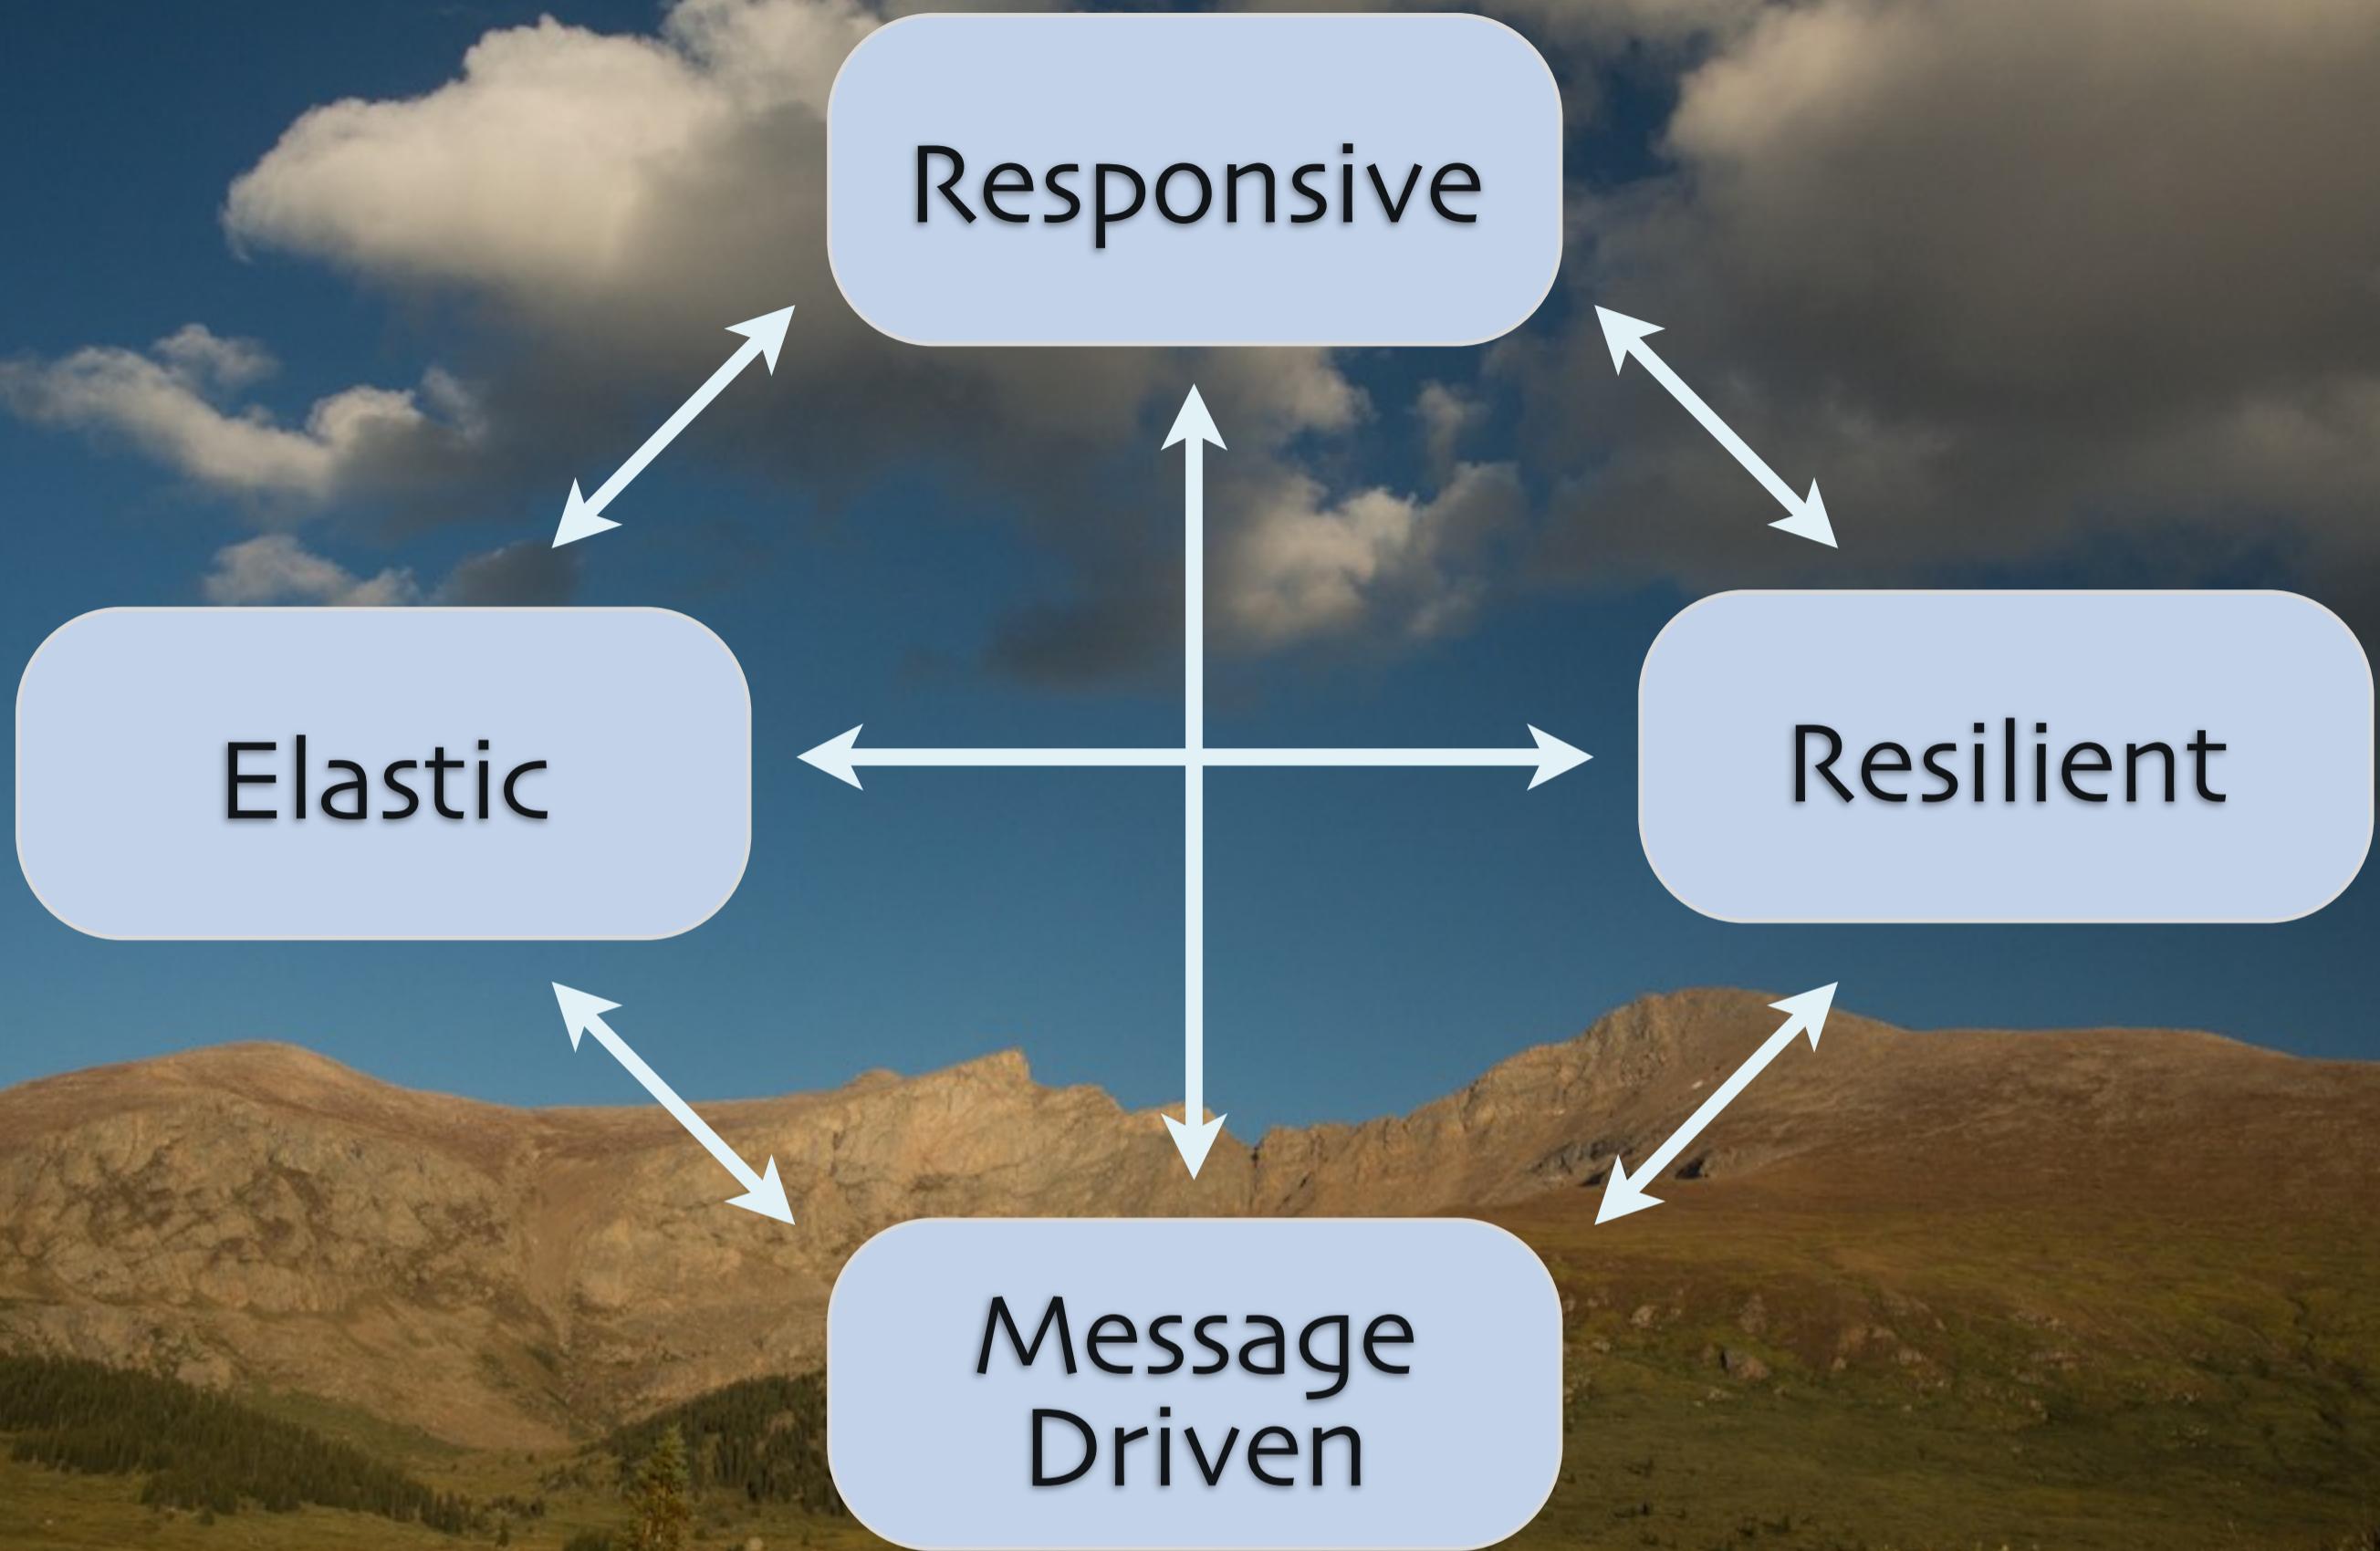


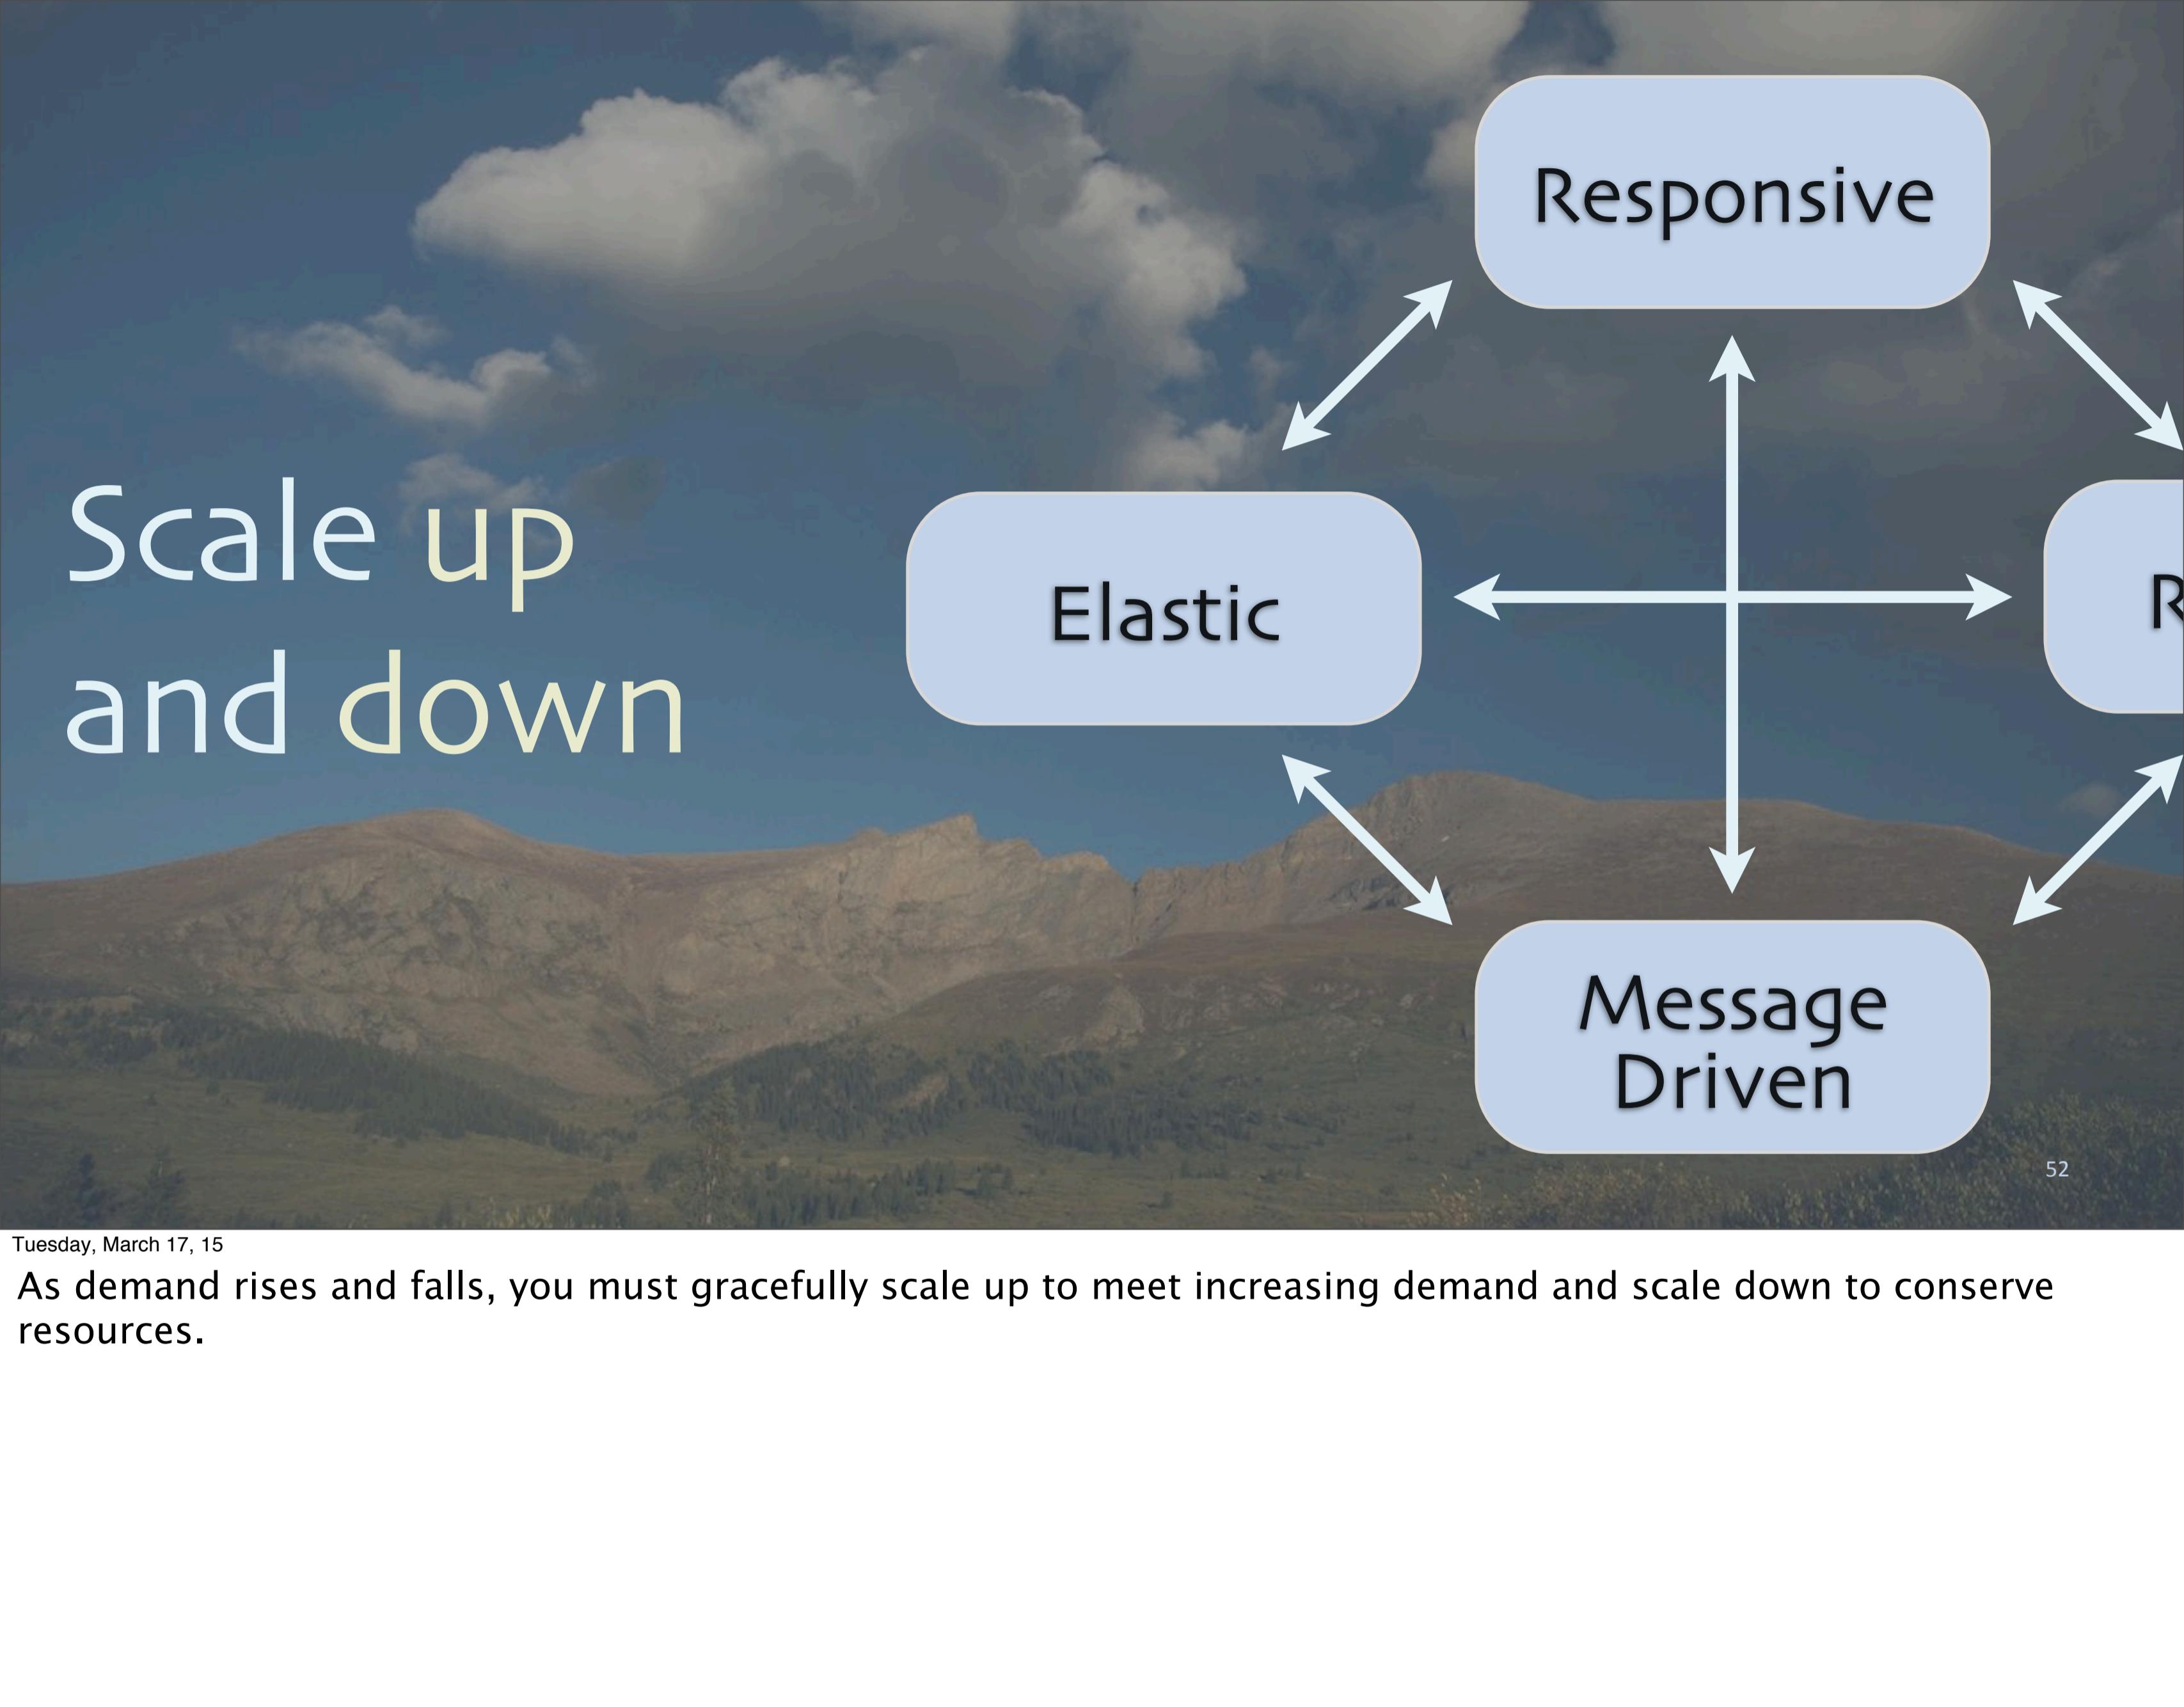
Clean separation
of normal processing
from recovery.



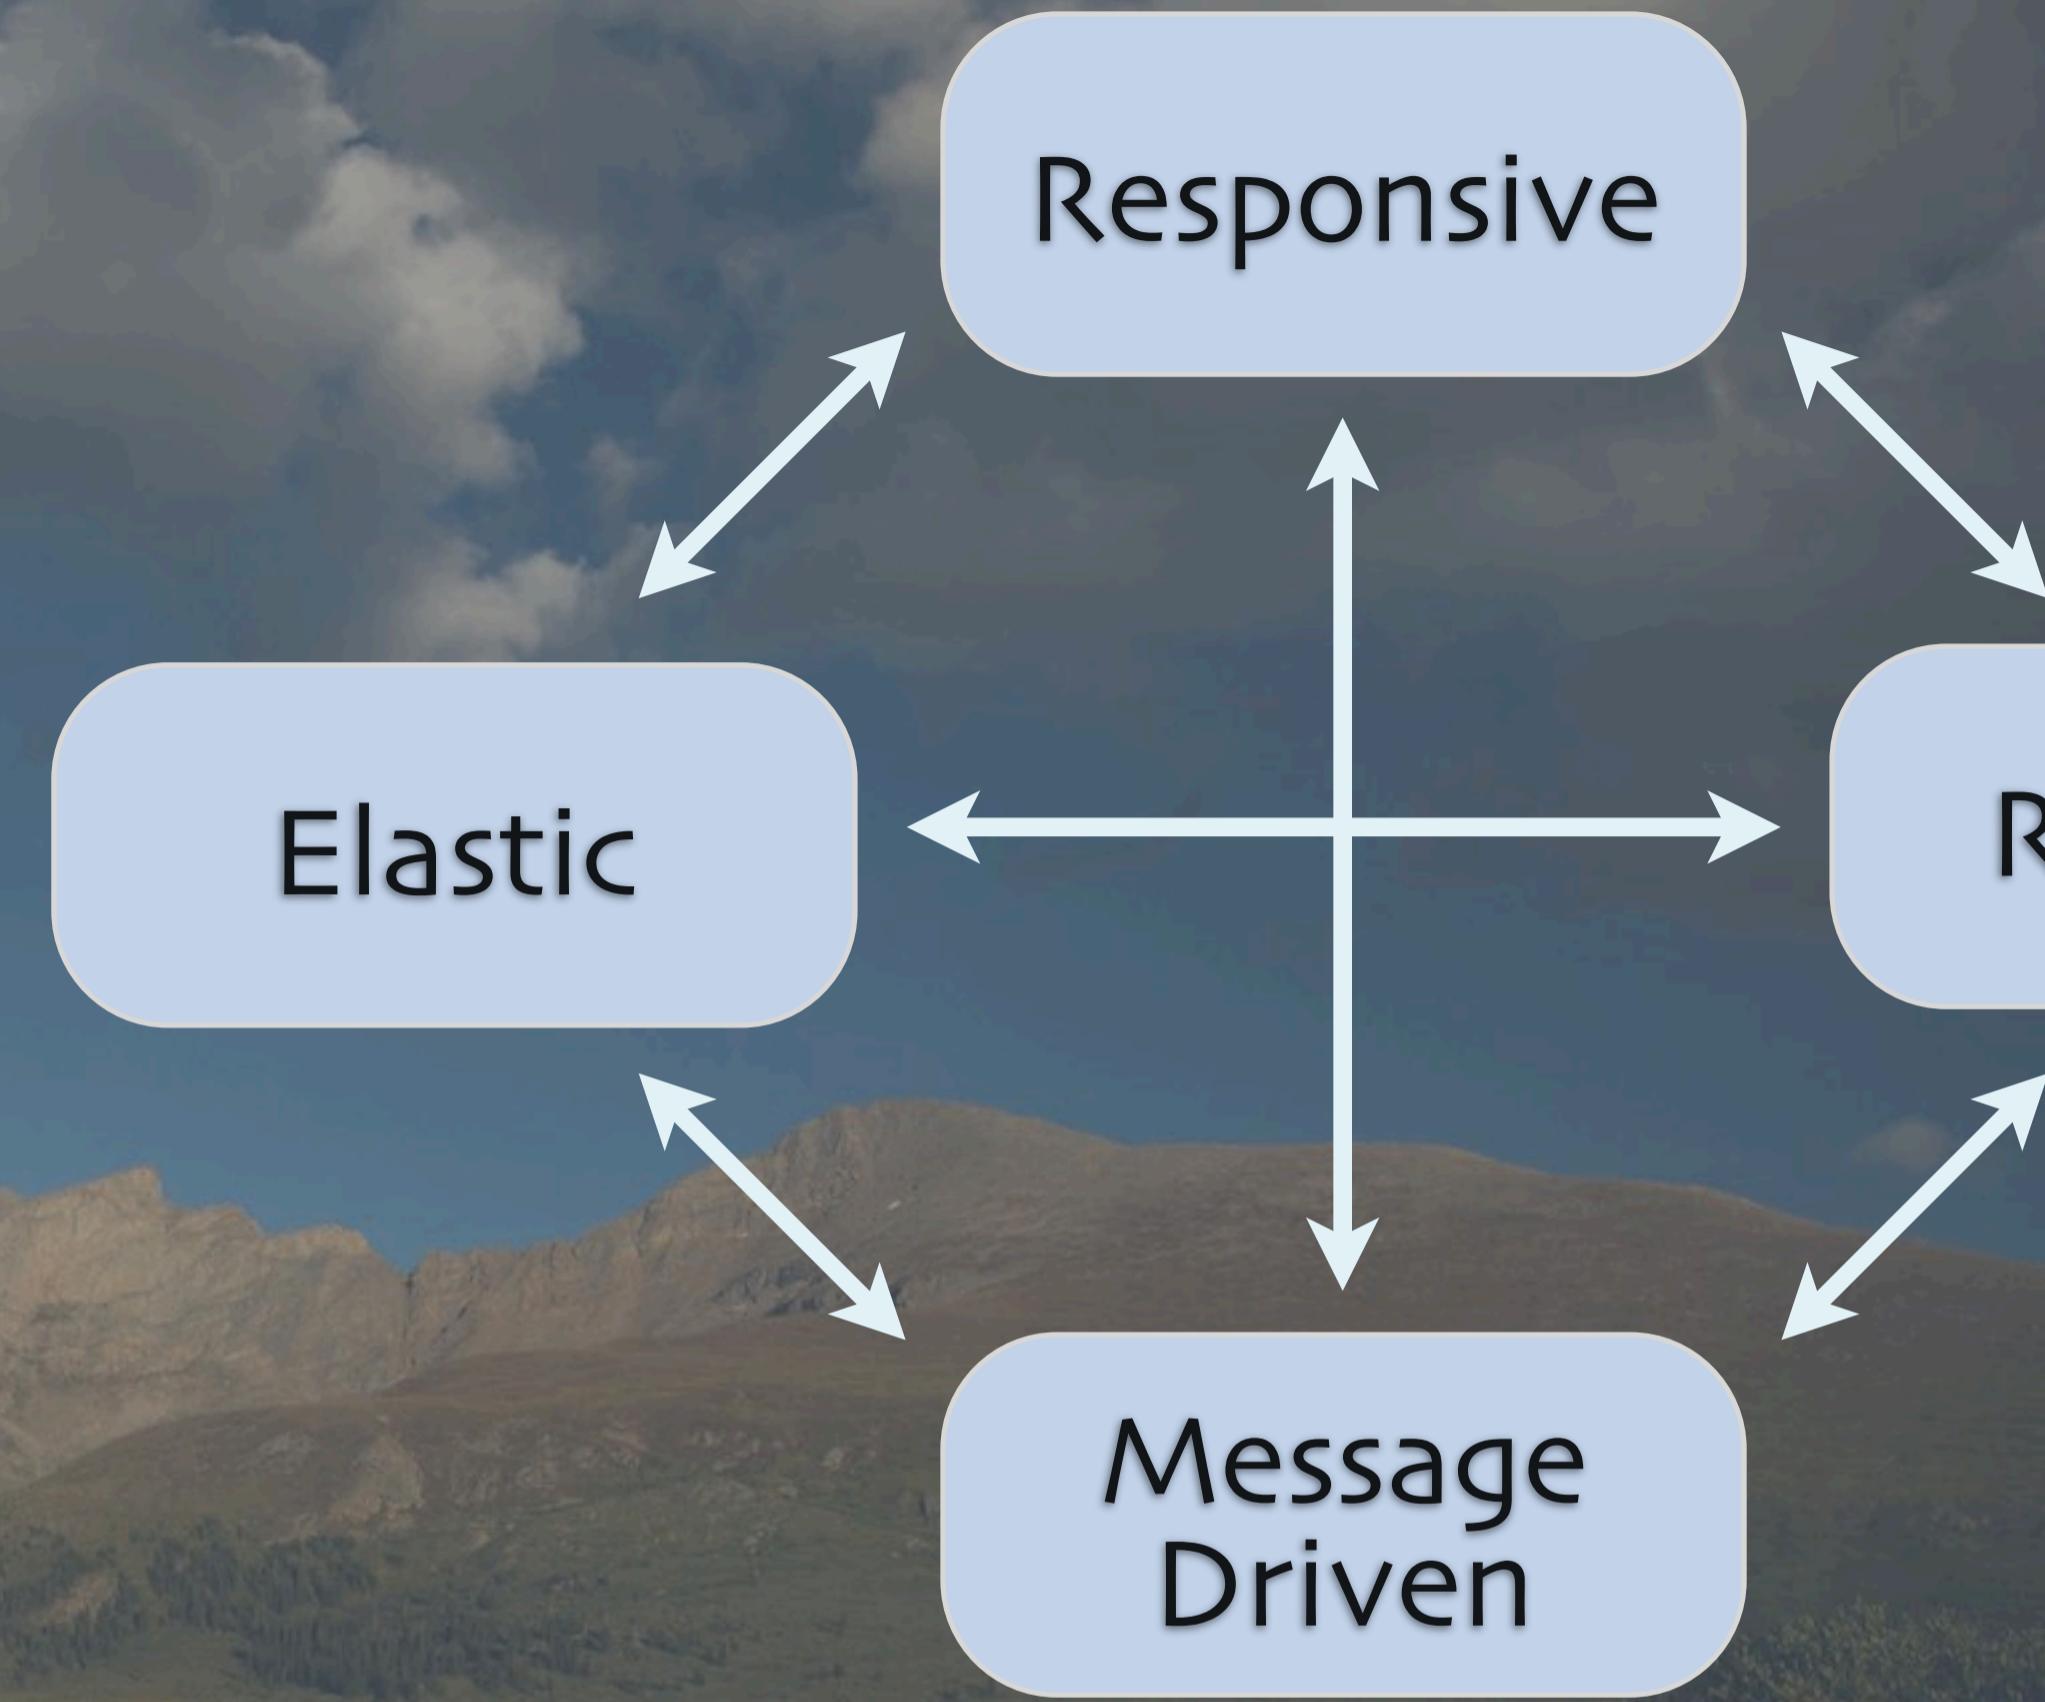
Not using Actors?
See **Hystrix**
from Netflix.

<https://github.com/Netflix/Hystrix>





Scale up and down



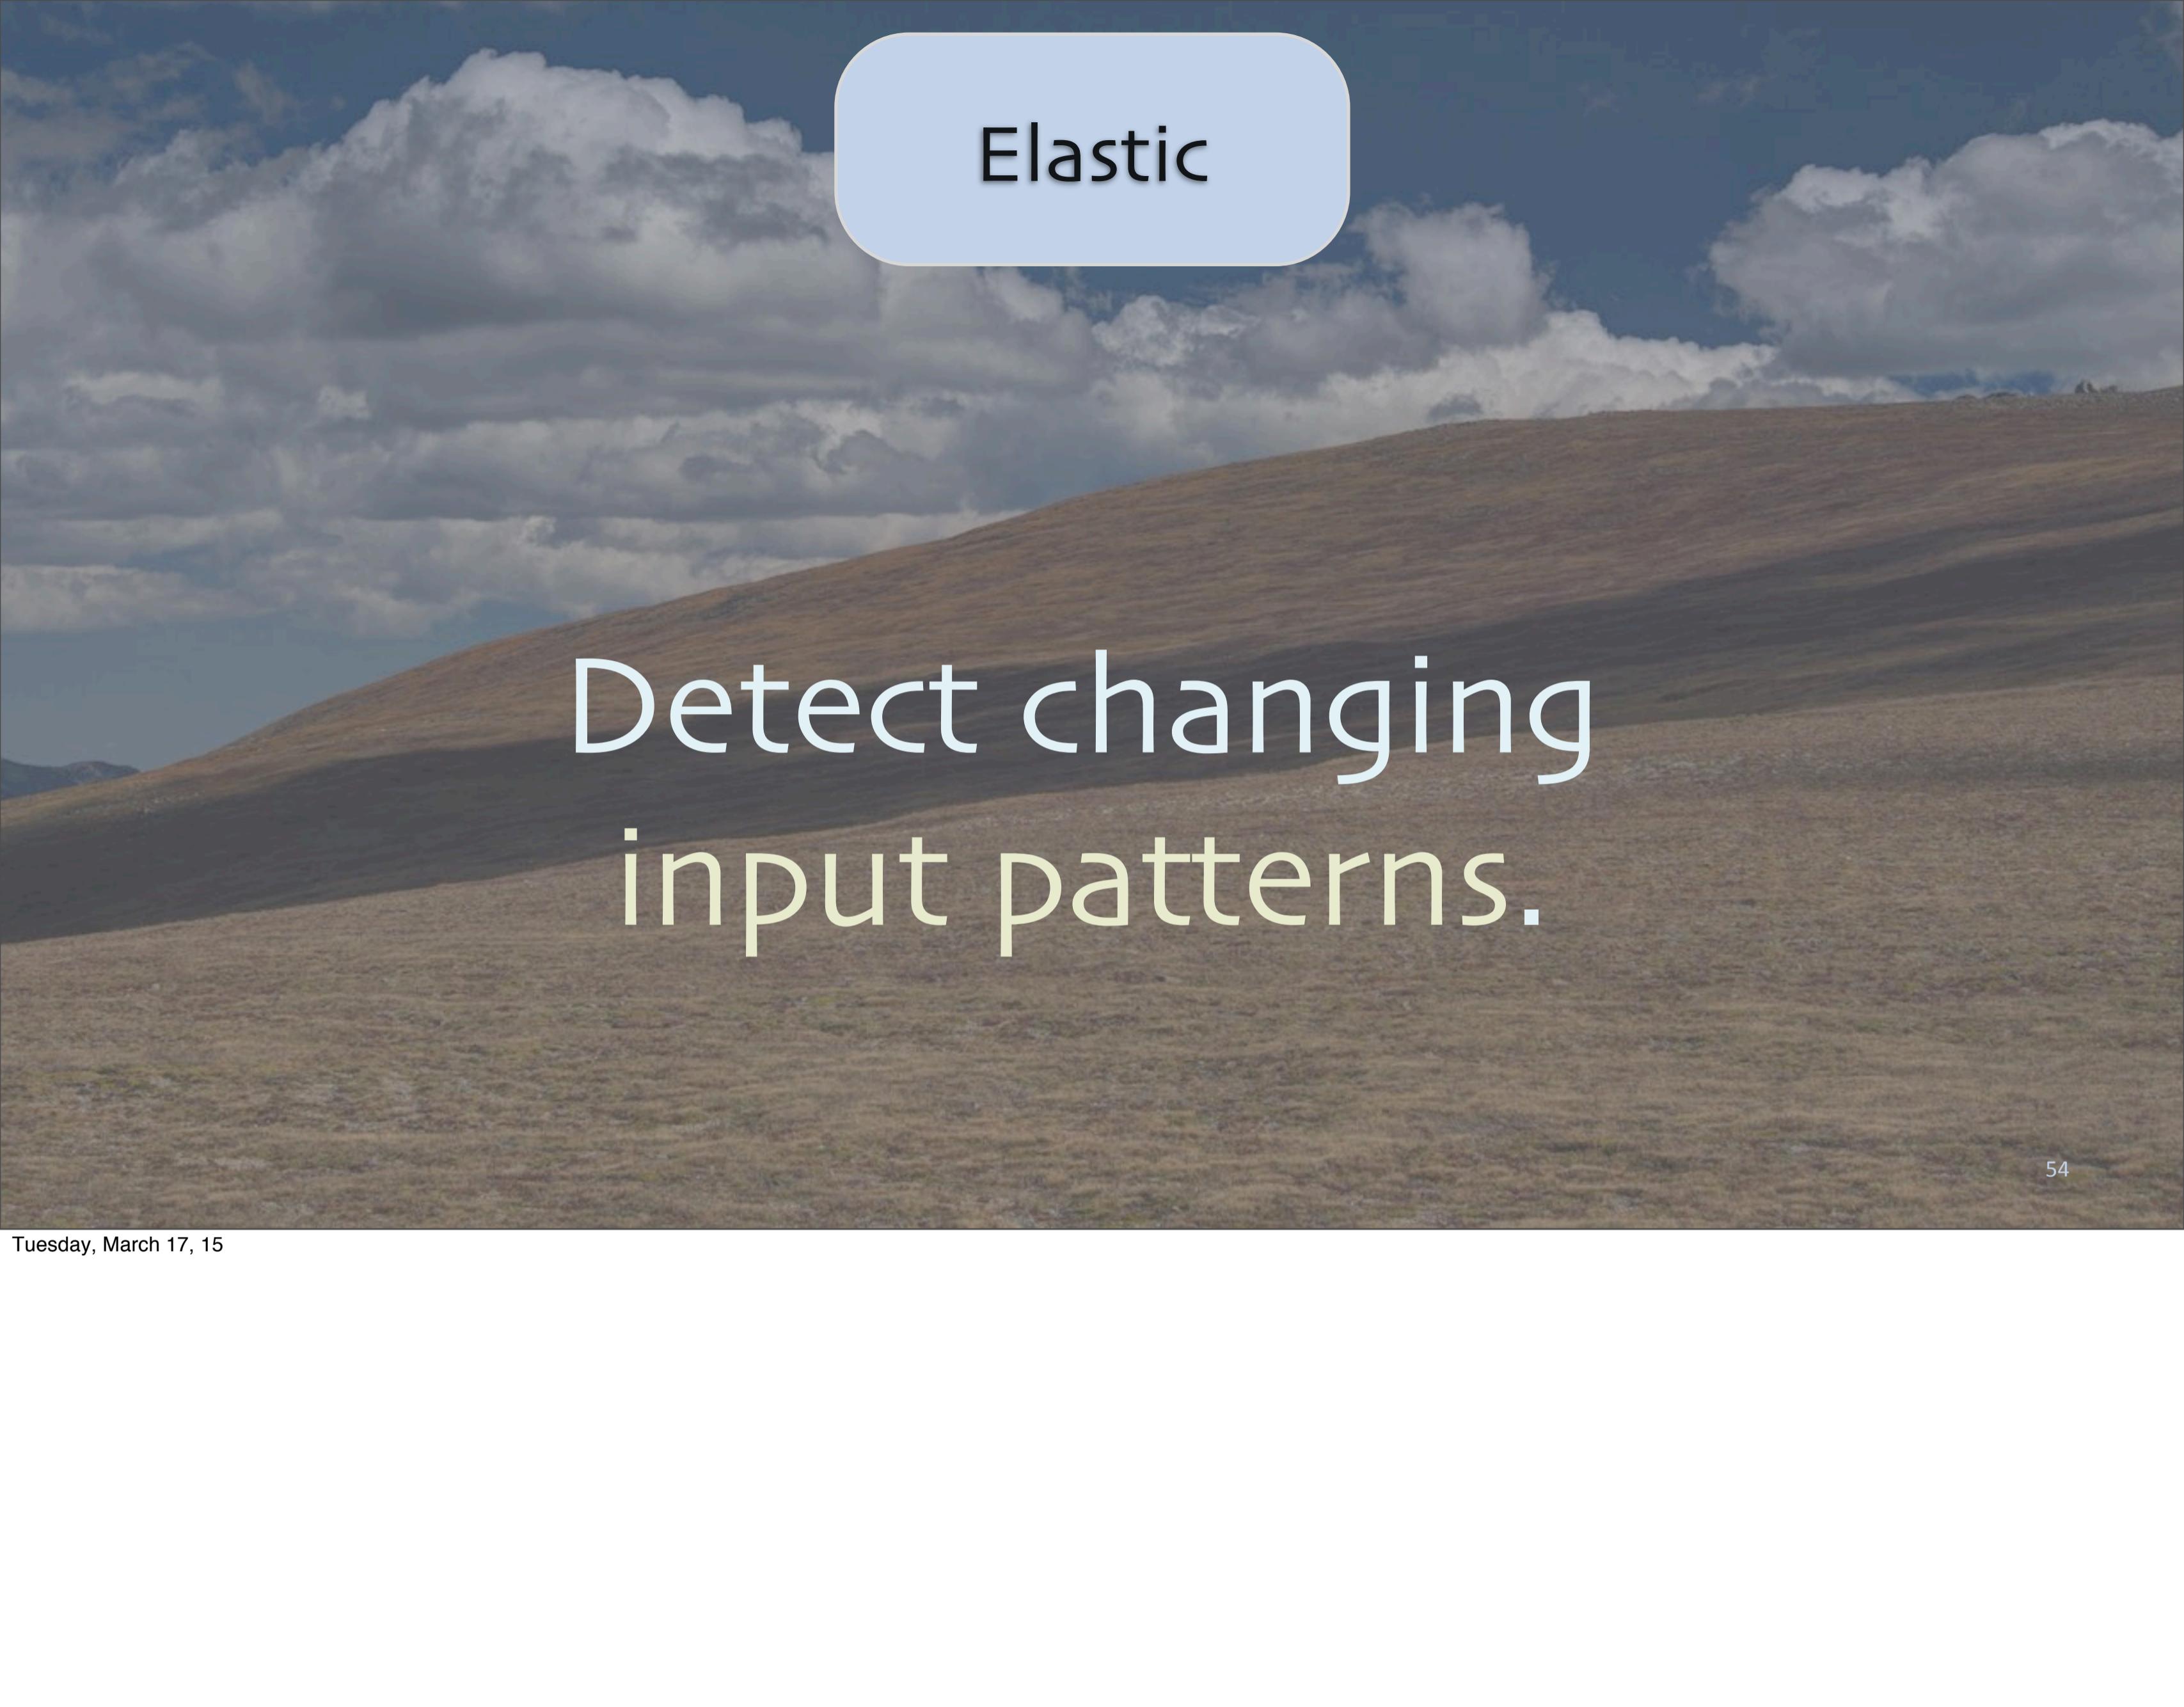
52

Tuesday, March 17, 15

As demand rises and falls, you must gracefully scale up to meet increasing demand and scale down to conserve resources.

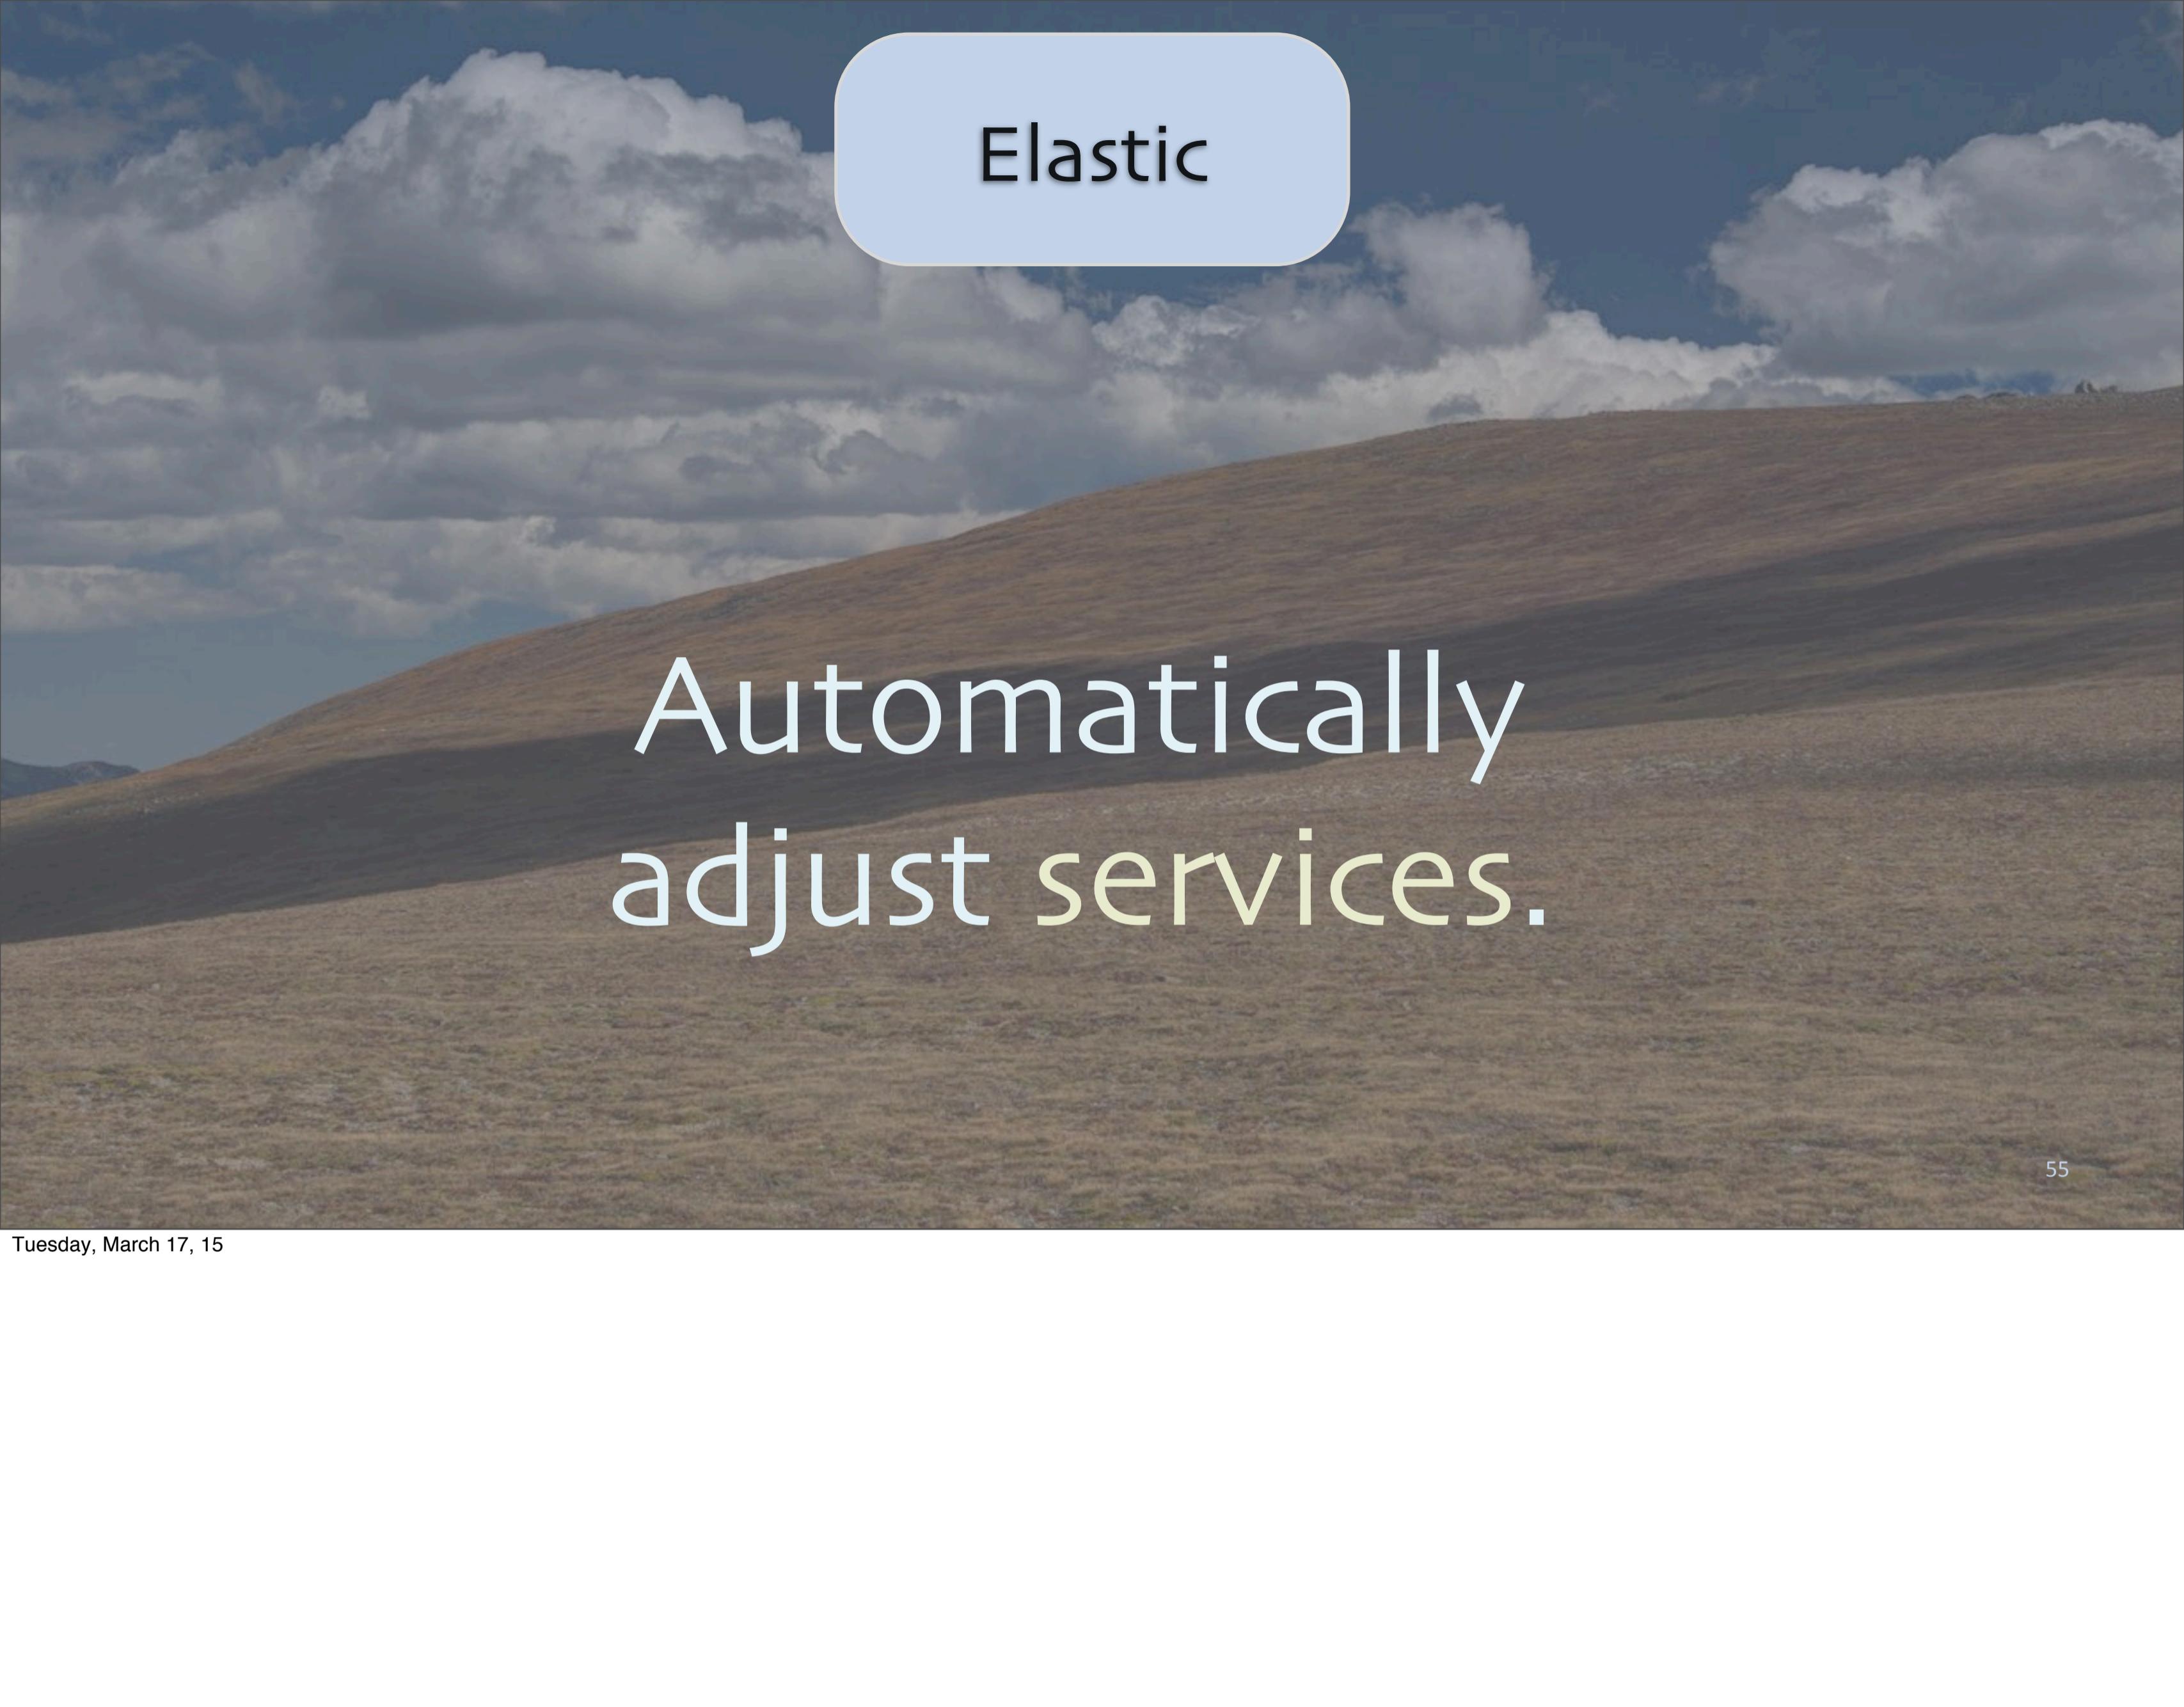


Elastic



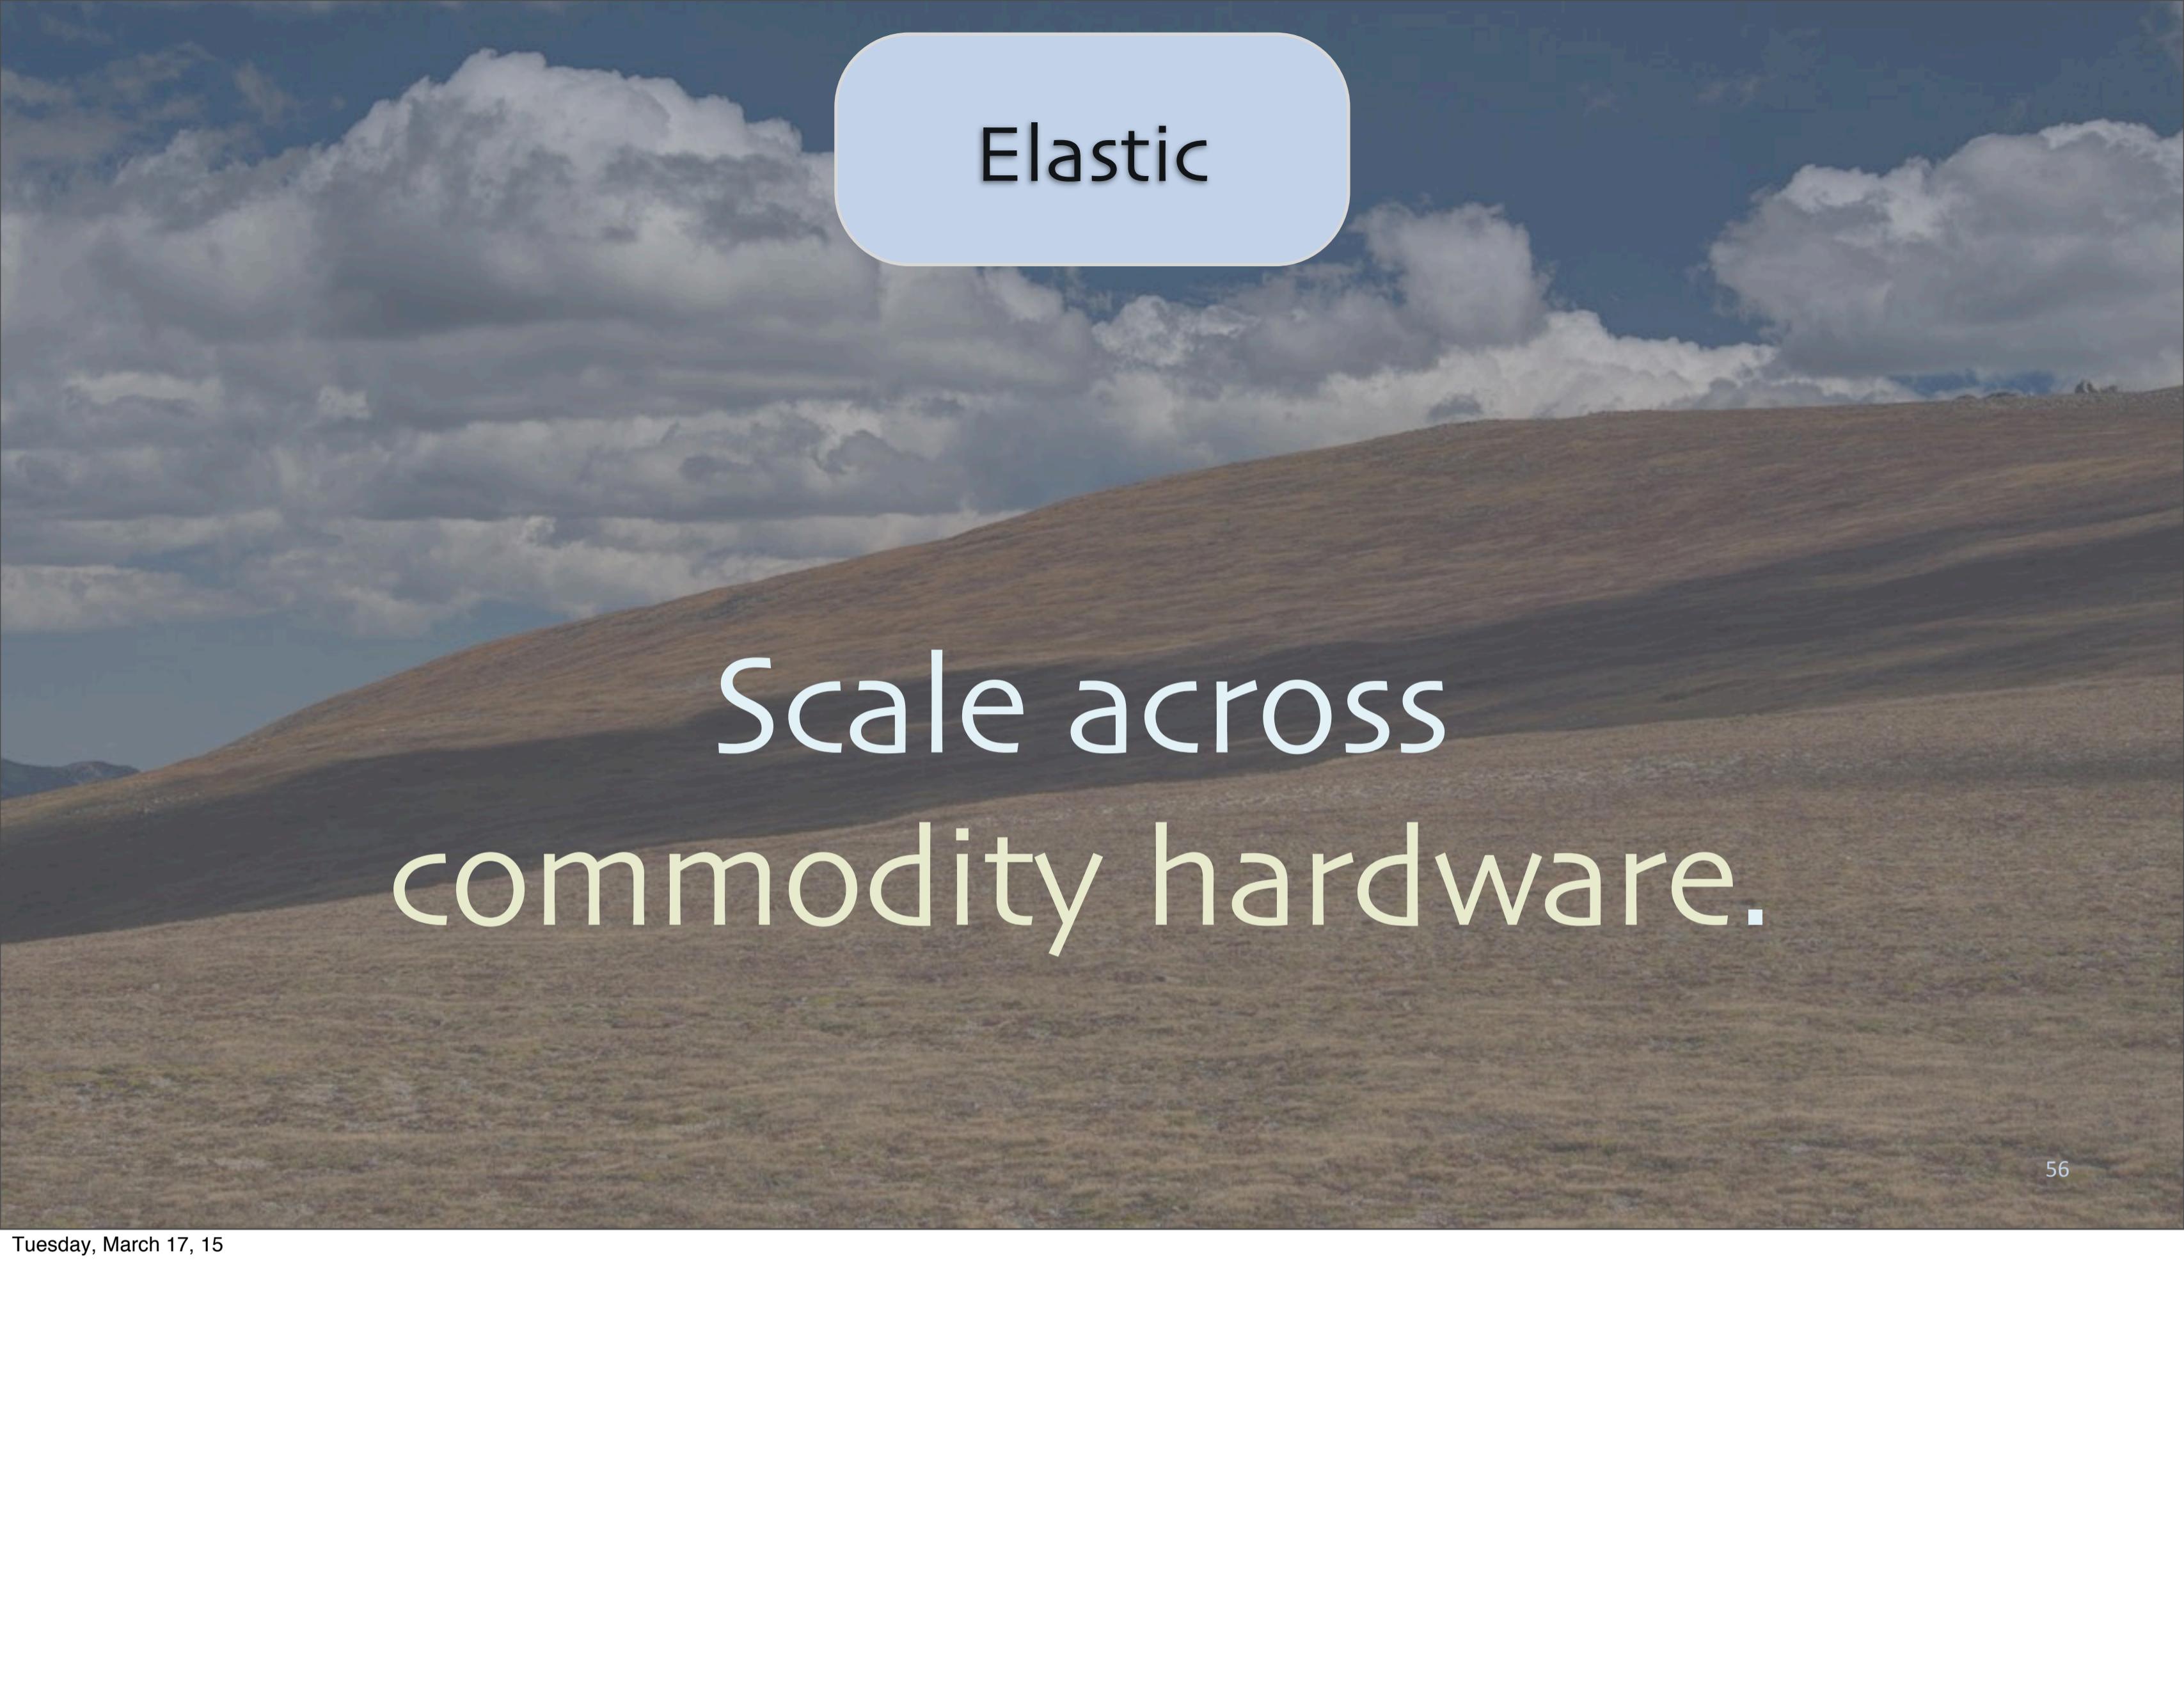
Elastic

Detect changing
input patterns.



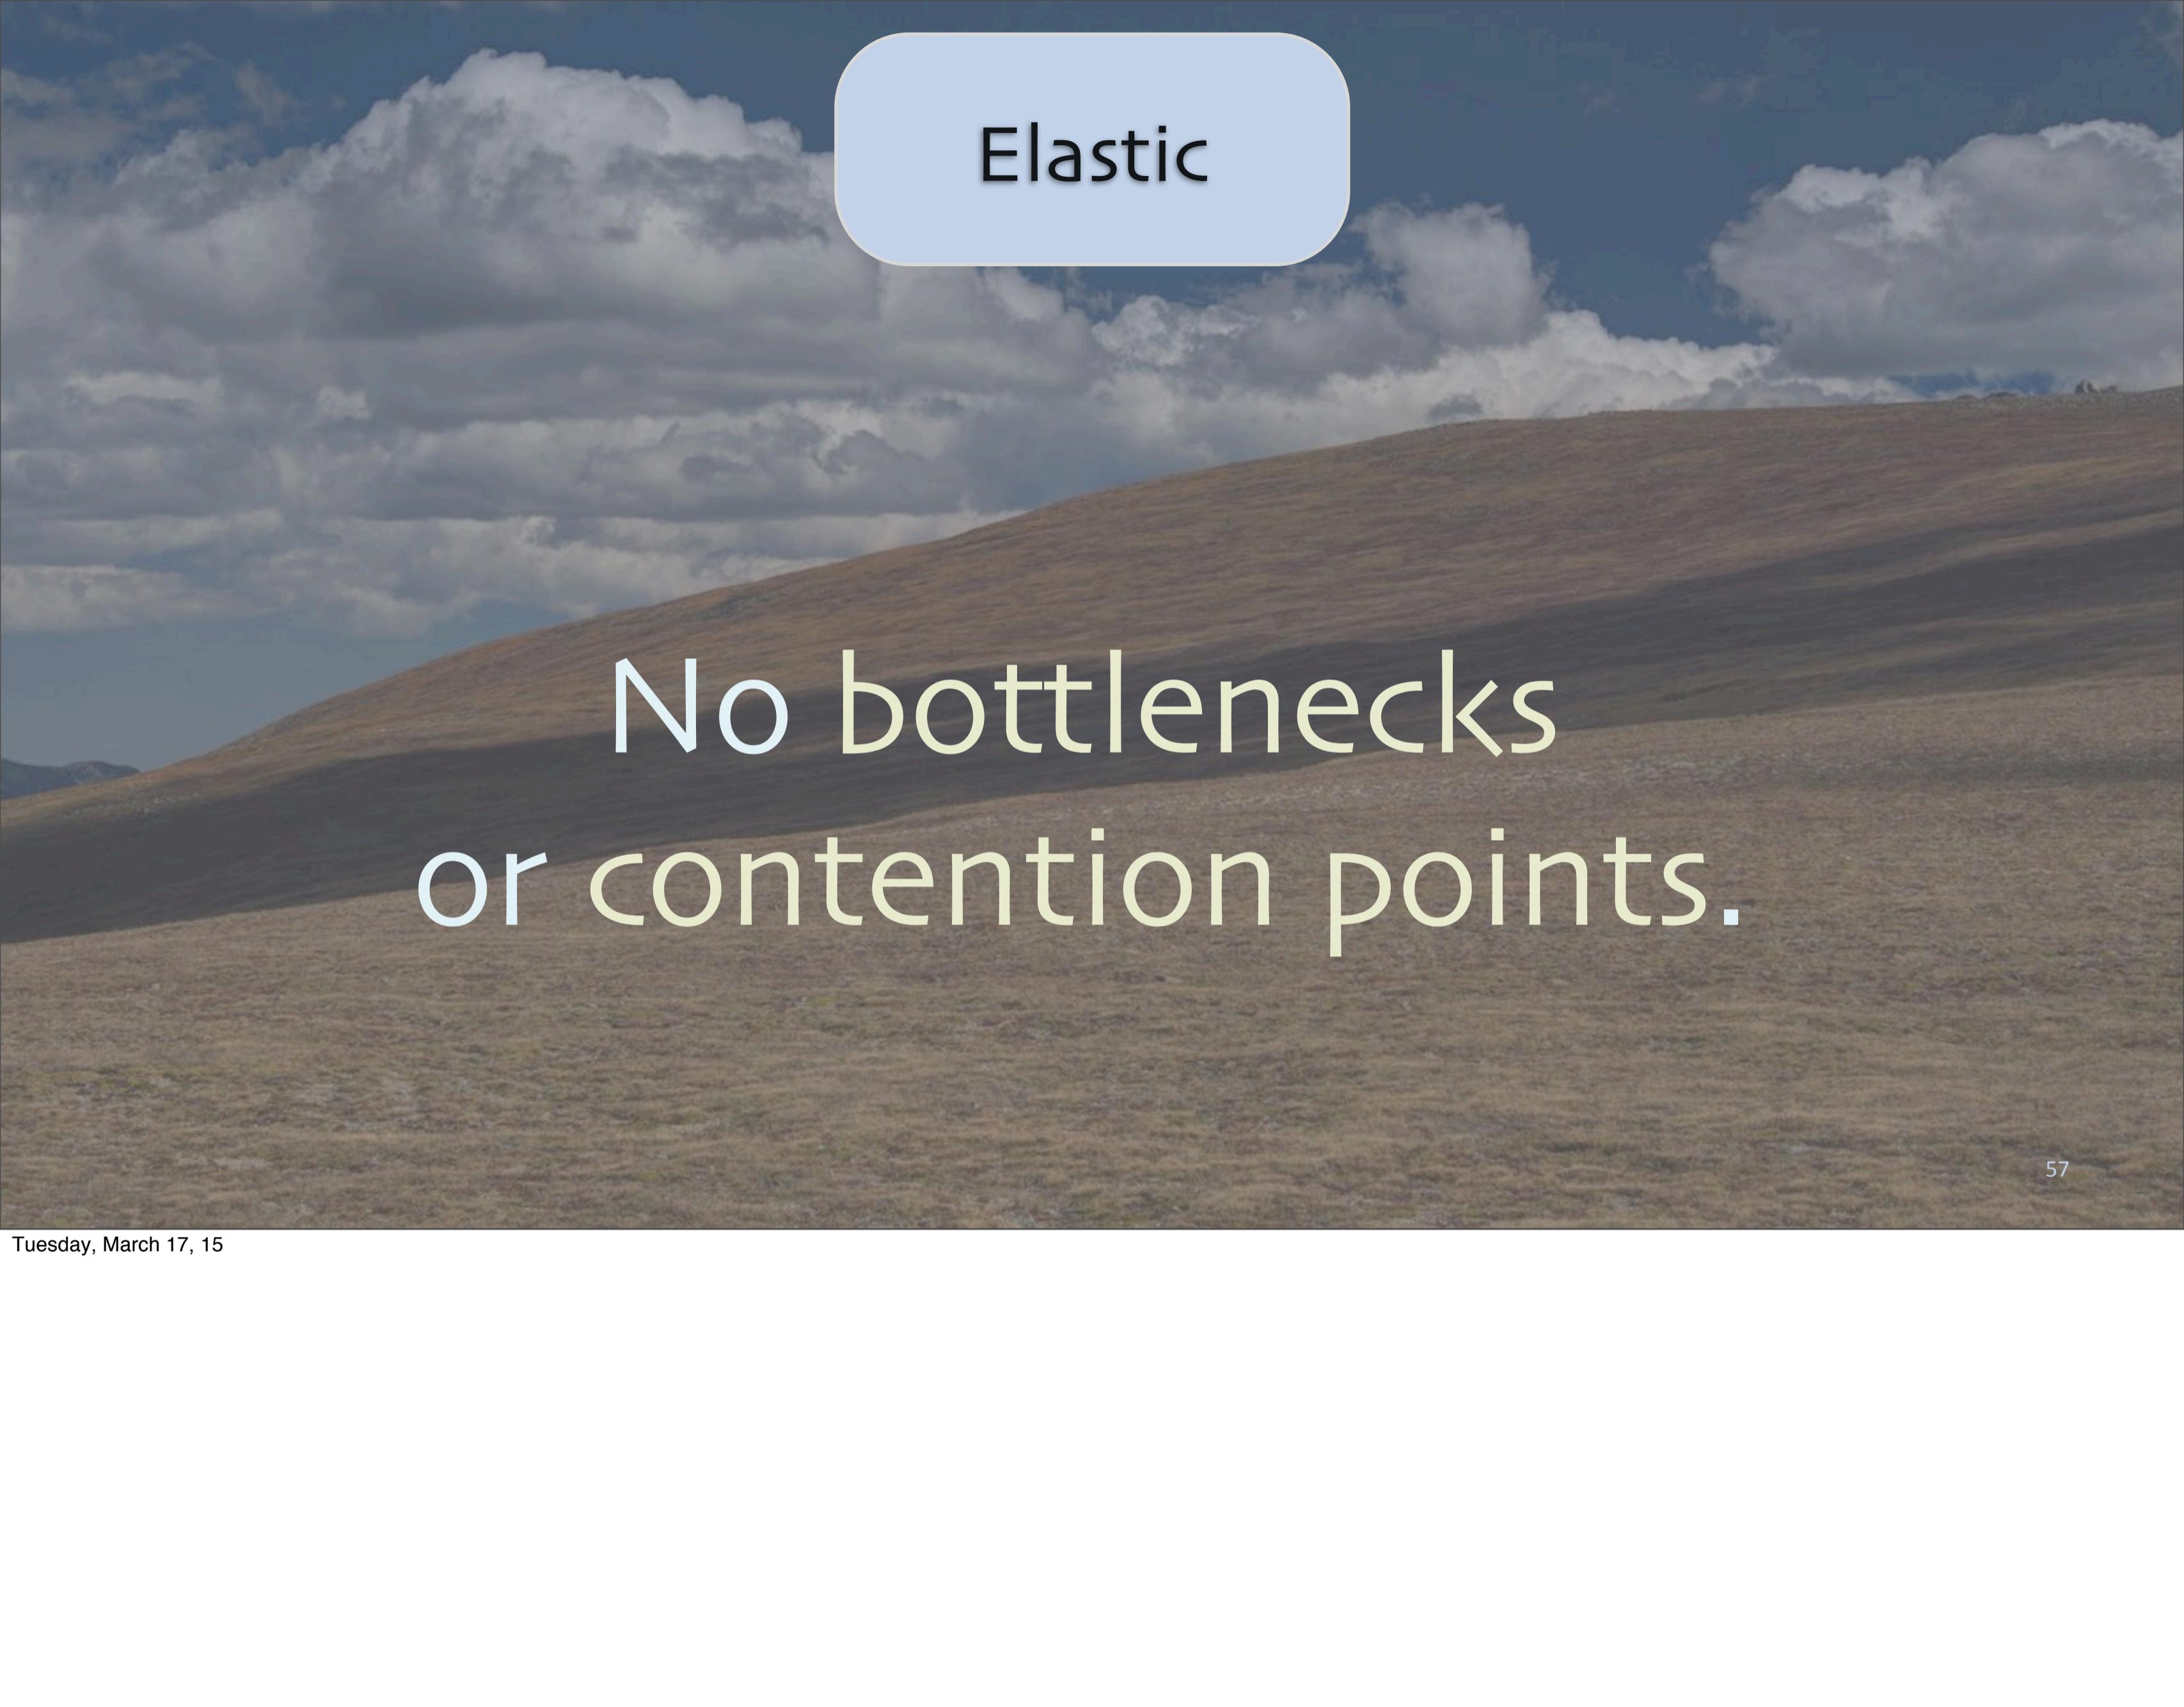
Elastic

Automatically
adjust services.



Elastic

Scale across
commodity hardware.

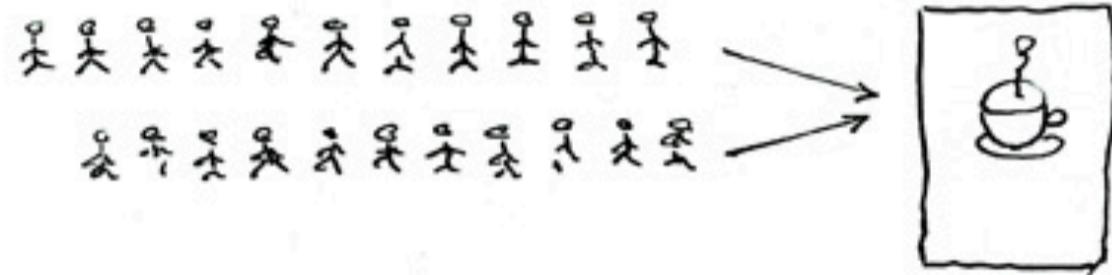
The background image shows a wide landscape with rolling, brownish-green hills under a sky filled with large, white, billowing clouds.

Elastic

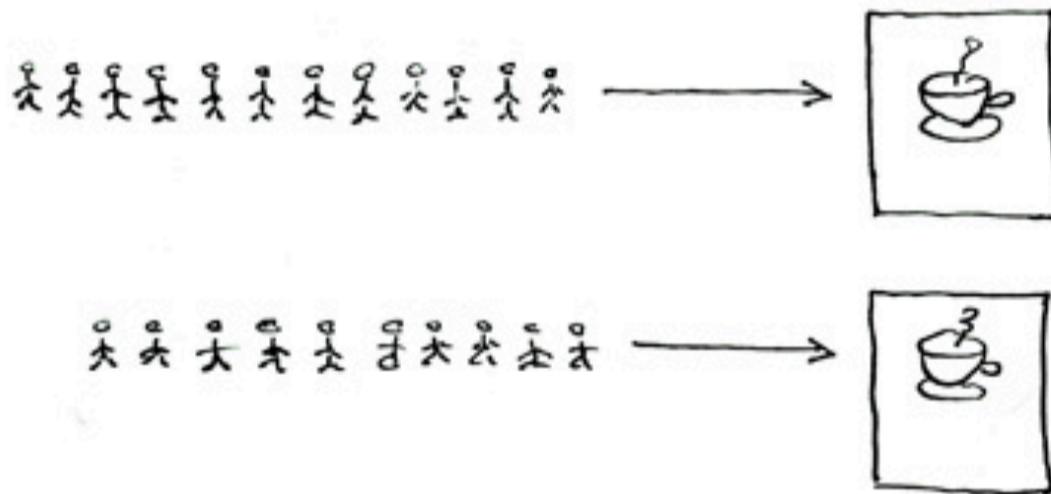
No bottlenecks
or contention points.

Elastic

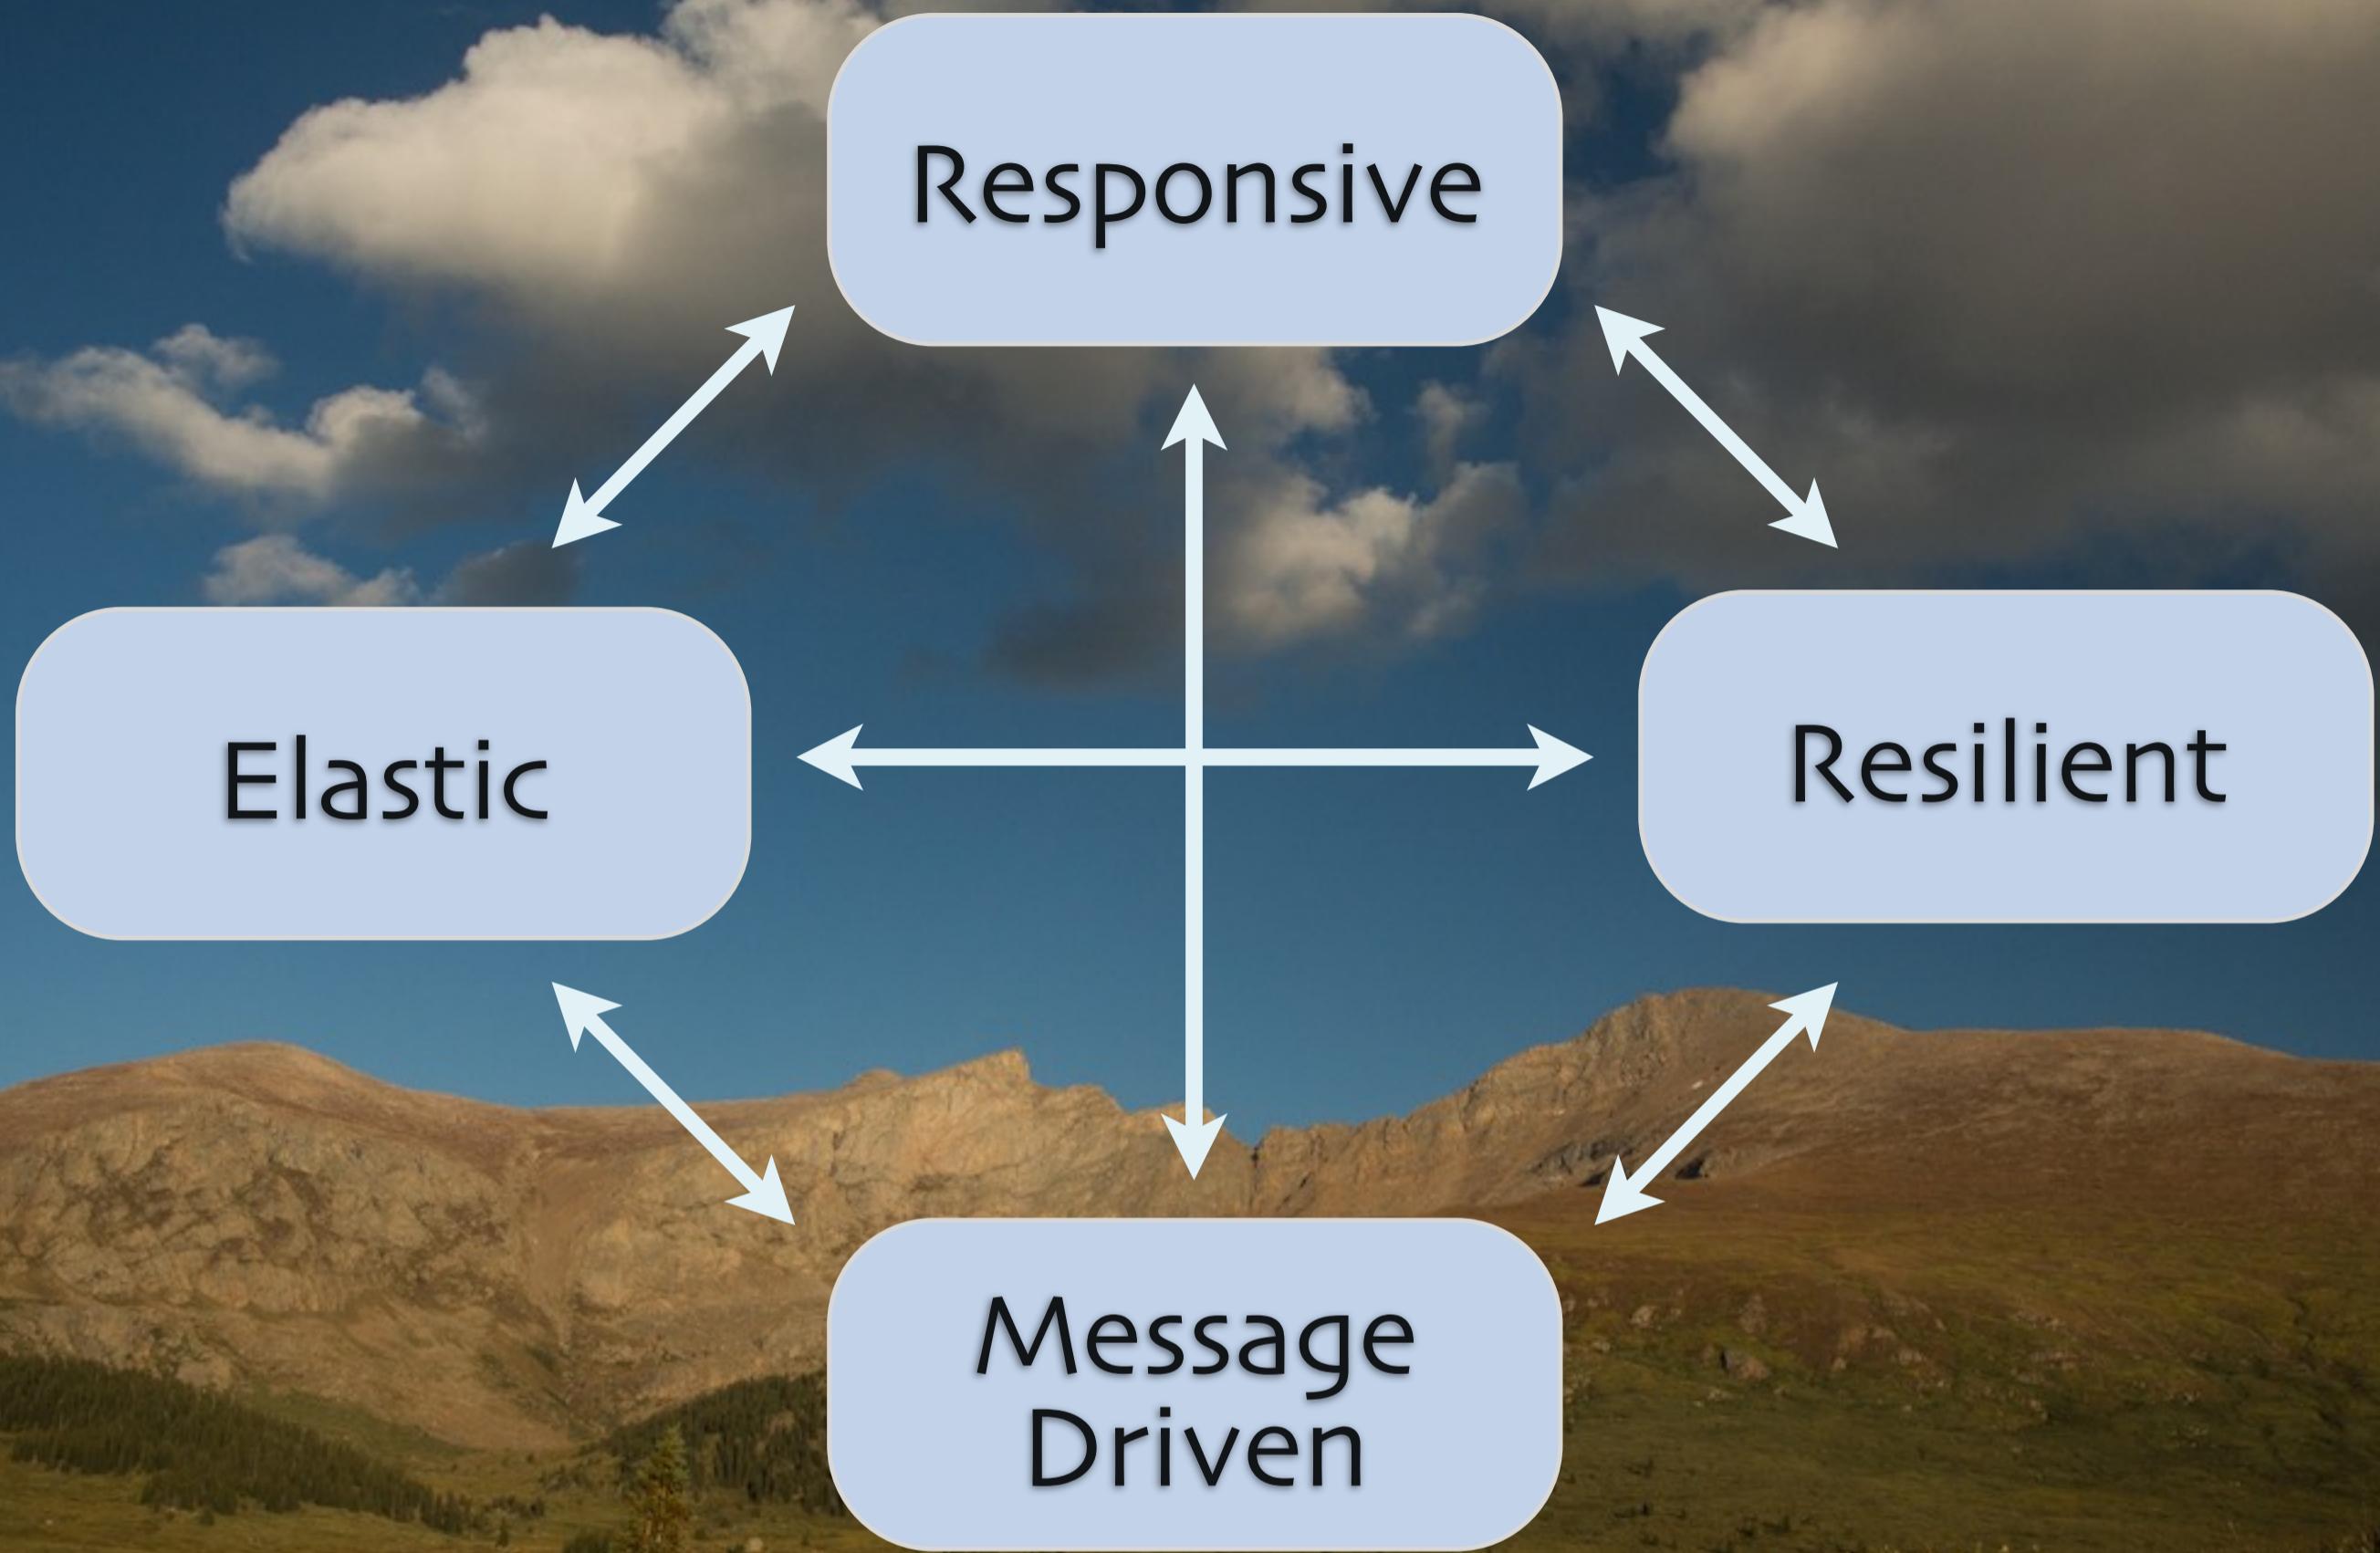
Concurrent = Two Queues One Coffee Machine



Parallel = Two Queues Two Coffee Machines



© Joe Armstrong 2013



```
graph TD; A([Elastic]) <--> B([Resilient]); A <--> C([Message Driven]); B <--> C
```

Elastic

Resilient

Message
Driven

To react, you must
be message driven.

60

Tuesday, March 17, 15

To be responsive to the world around you, you must interact with the world through messages.



Message
Driven



Message
Driven

Asynchronous
message passing.

62



Message Driven

Defines boundaries,
promotes loose coupling
and isolation.



Message
Driven

Promotes location transparency.
Handle errors as messages.



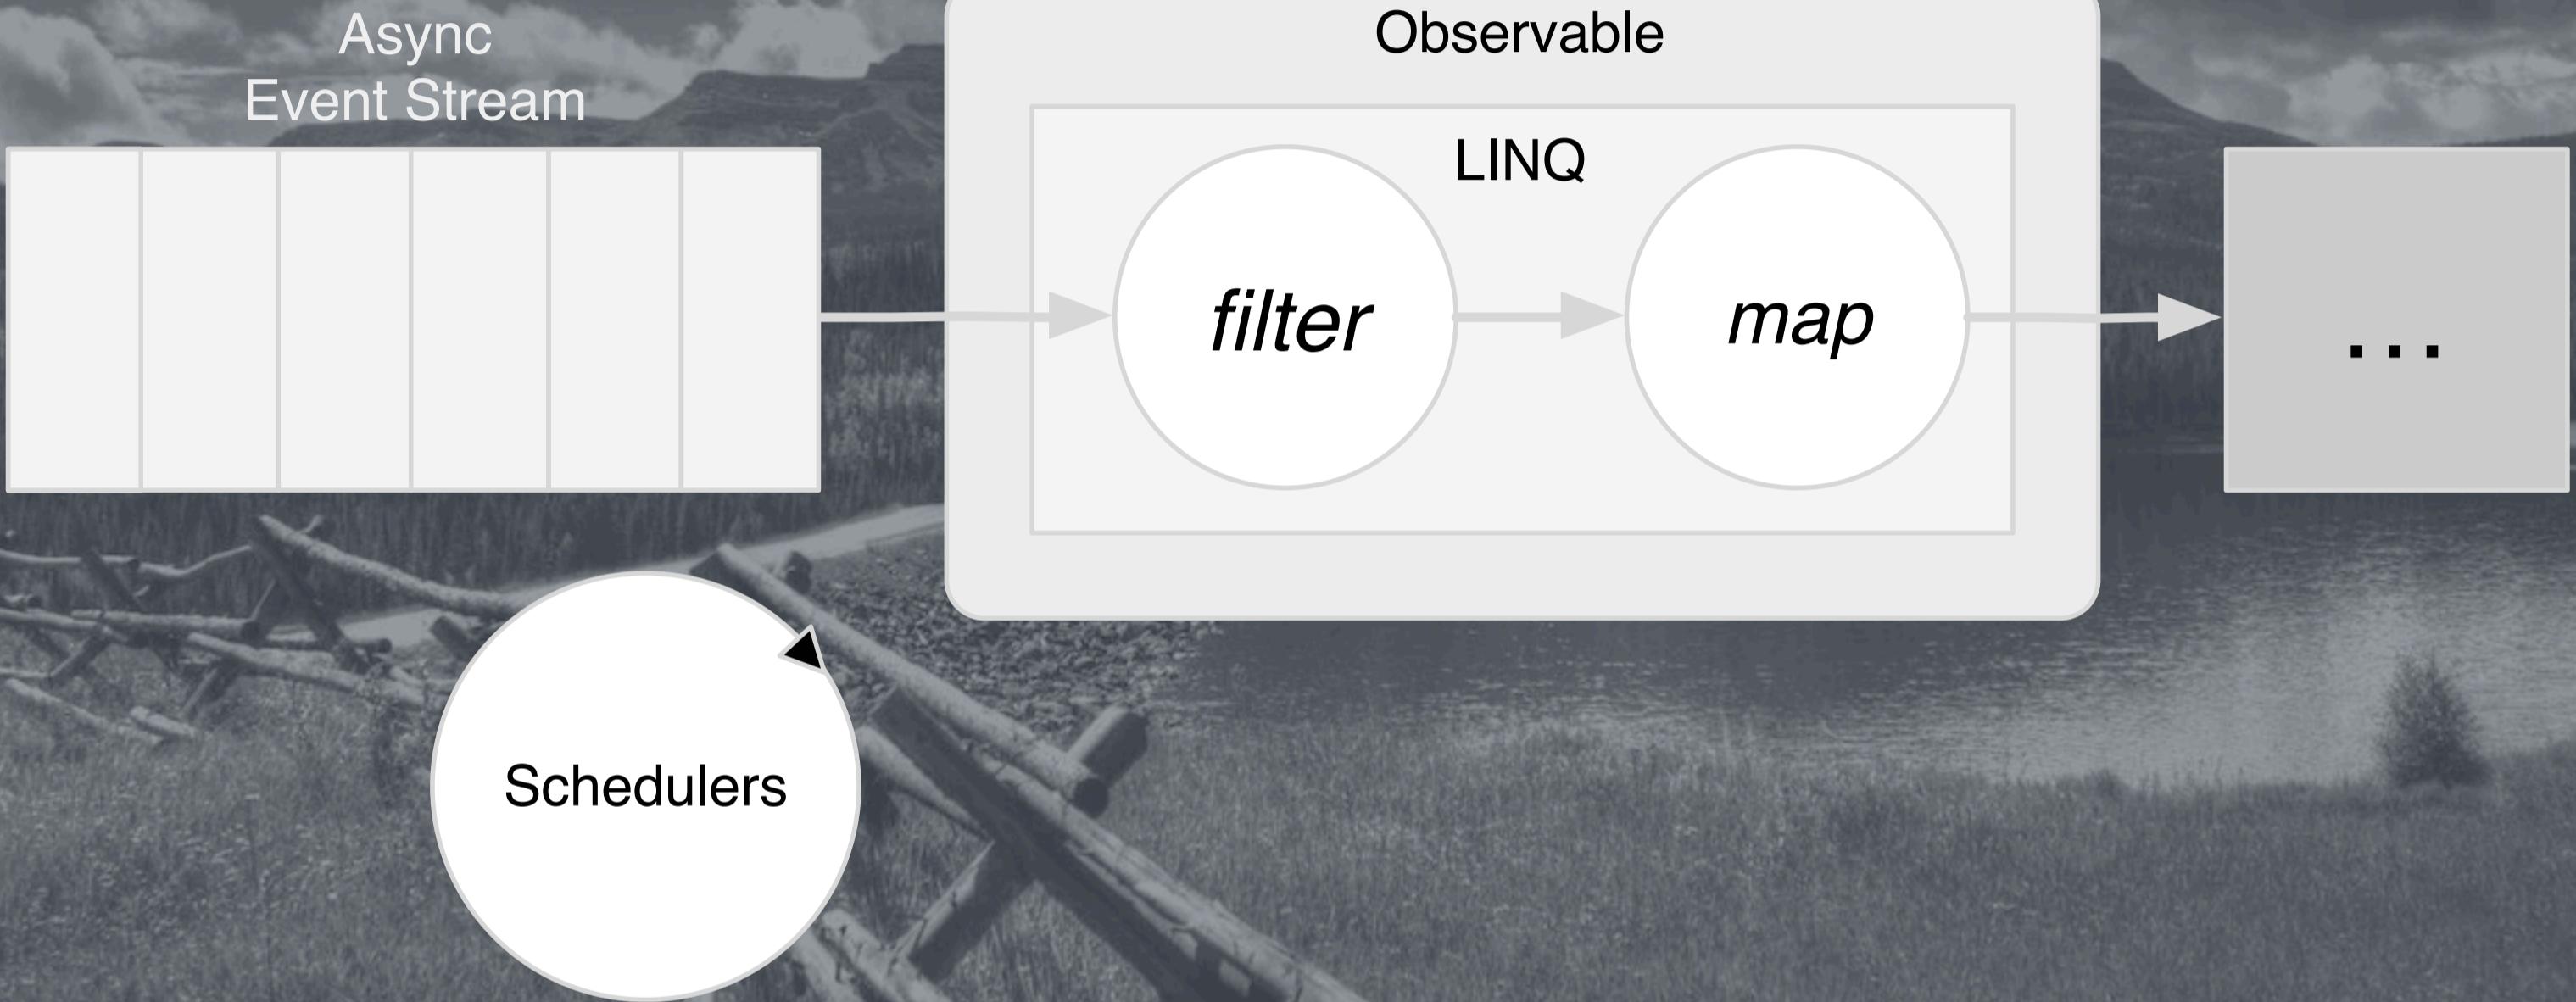
Message
Driven

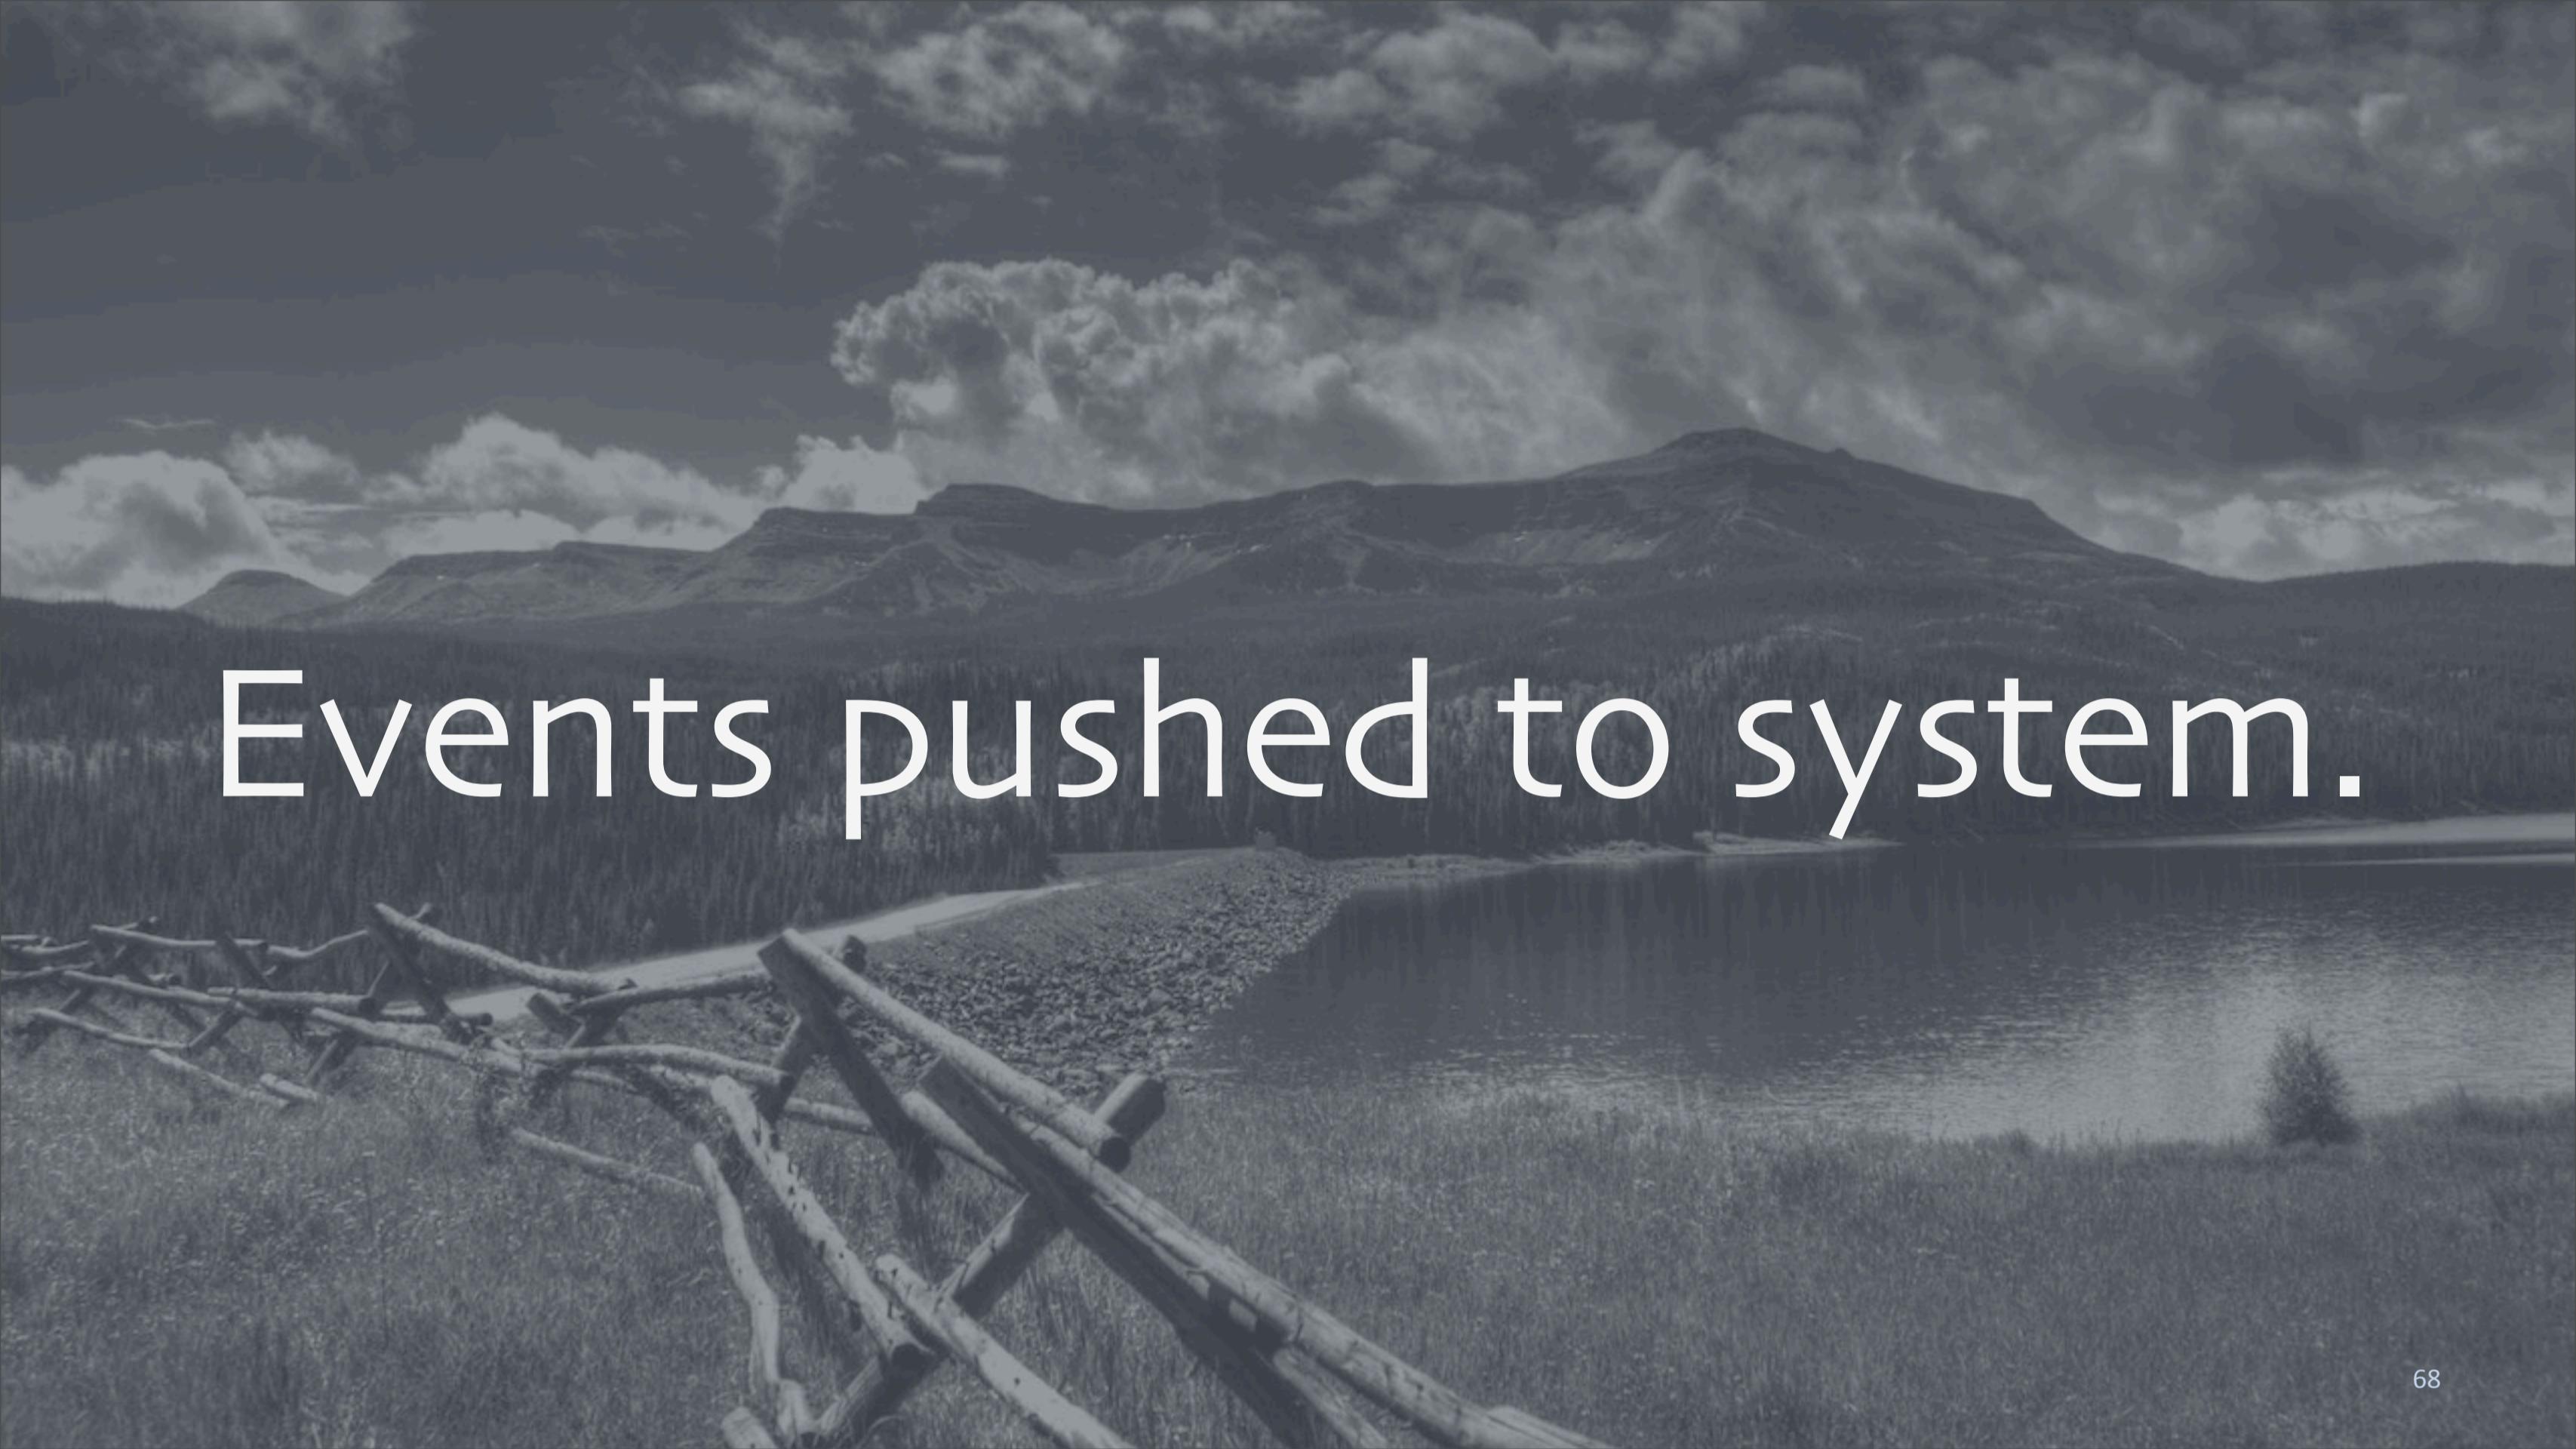
Promotes global management
and flow control
through backpressure.

65

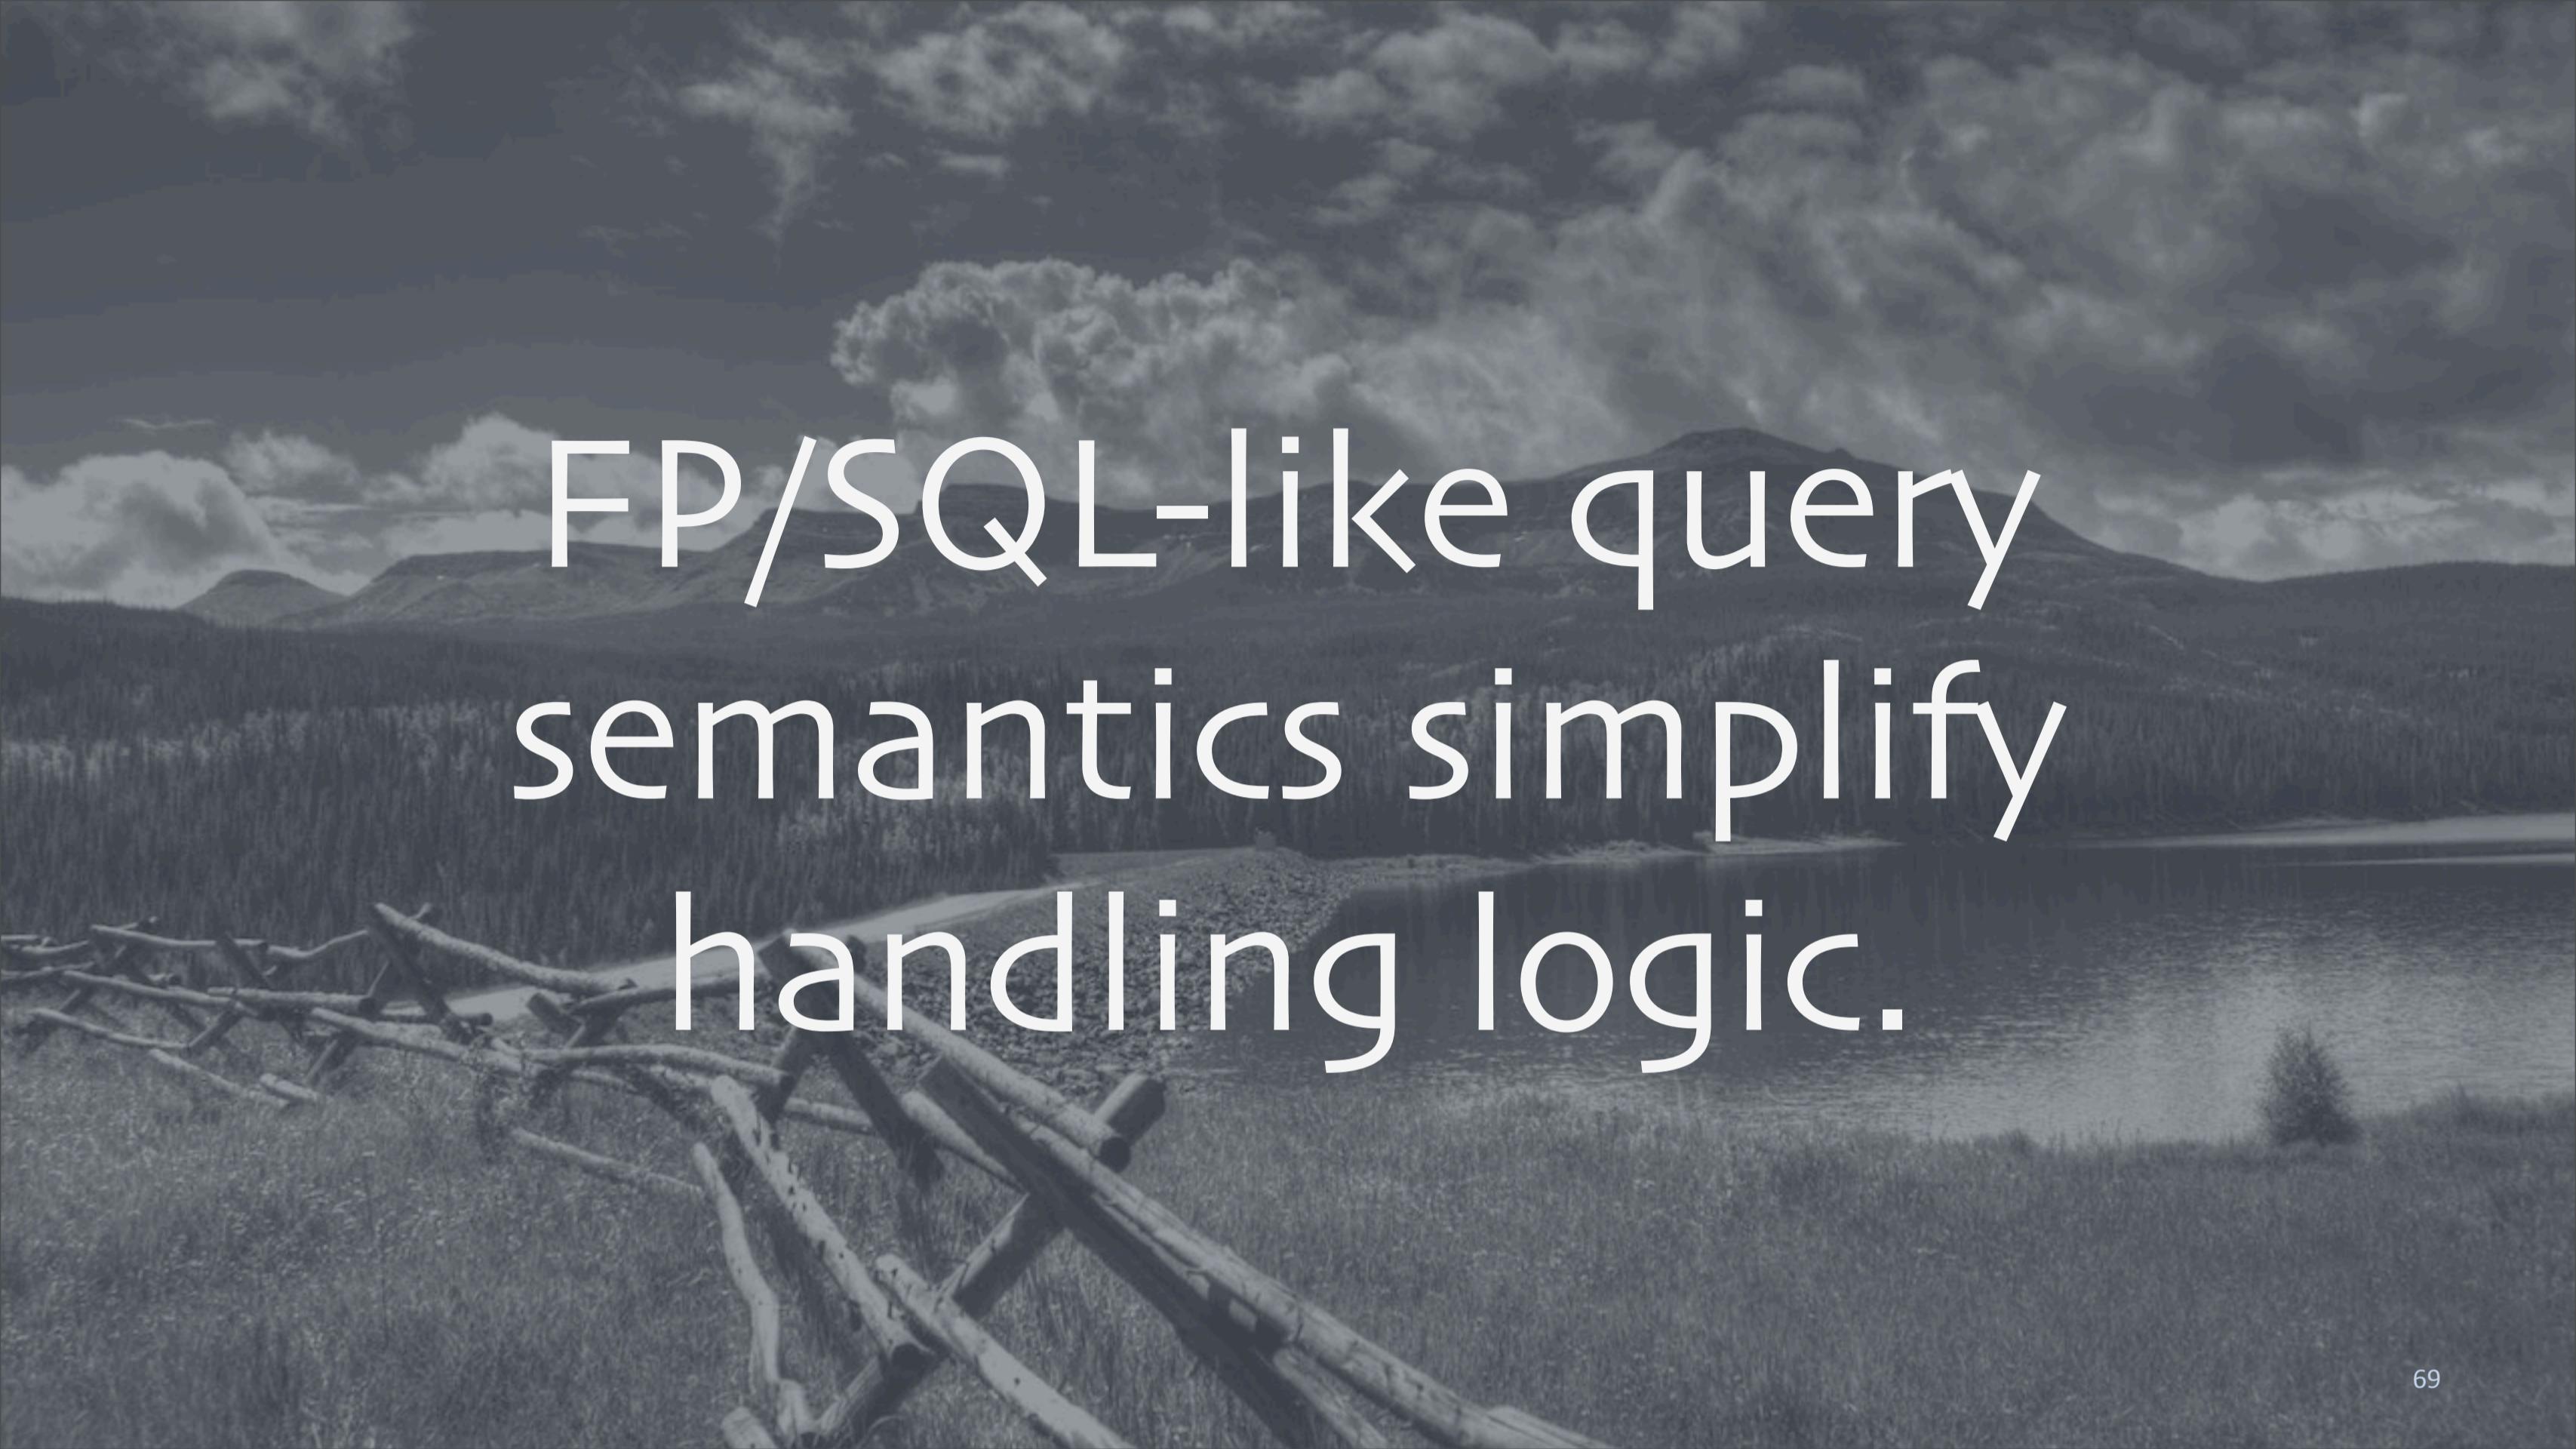
Reactive Extensions (Rx)

LINQ Rx





Events pushed to system.



FP/SQL-like query
semantics simplify
handling logic.

How: Reactive Tools



70

Tuesday, March 17, 15

We mentioned a few already, let's fill in some details. This won't be an exhaustive list.
Hat tip to Jamie Allen at Typesafe for some of these ideas.

How: Reactive Tools

- Functional Programming
- Distributed Computing “Laws”
- Software Transactional Mem.
- Event Loops

71

Tuesday, March 17, 15

We mentioned a few already, let's fill in some details. This won't be an exhaustive list.

How: Reactive Tools

- CSP
- Futures
- Actors - Erlang or Akka
- Rx and variants
- Reactive Streams

72

Tuesday, March 17, 15

CSP – Communicating Sequential Processes.

Functional Programming



73

Tuesday, March 17, 15

Functional Programming

Need to be asynchronous
and nonblocking.
Avoid locks.

Functional Programming

Prefer immutable values and
side-effect free functions.

Functional Programming

Objects - suitable for modules.
Functional - for everything else.

Architecture Side Note:

The biggest mistake of OOP
was the idea that we should
faithfully model the world
in code.

Distributed Computing



78

Tuesday, March 17, 15

Distributed Computing

Serializability (order) and
Linearizability (change history
results in same order?).
CRDTs, Lattices.

79

Tuesday, March 17, 15

CRDTs – Commutative Replicated Data Types (<http://pagesperso-systeme.lip6.fr/Marc.Shapiro/papers/RR-6956.pdf>)
Lattices are more general concept applied here.

Software Transactional Memory



80

Tuesday, March 17, 15

Popularized first in Hardware, then in Software by Haskell. Now used in persistent datastructures.

Software Transactional Memory

Principled local state
mutation through transactions.
Limited scalability,
no distribution.

81

Tuesday, March 17, 15

Great description of STM by Simon Peyton-Jones, from the O'Reilly Book Beautiful Code, <http://research.microsoft.com/en-us/um/people/simonpj/papers/stm/#beautiful>

Event Loops



82

Tuesday, March 17, 15

As explicitly managed in tools.

Event Loops

Pull from a thread.
Limited resiliency, scalability.
Callback hell...

83

Tuesday, March 17, 15

Many of these systems don't provide facilities for distribution or error recovery. Scalability can also be limited by the available threads, for those systems that pin a thread to each control flow construct. Some use a callback programming model exclusively, which quickly becomes unwieldy.

Communicating Sequential Processes

84

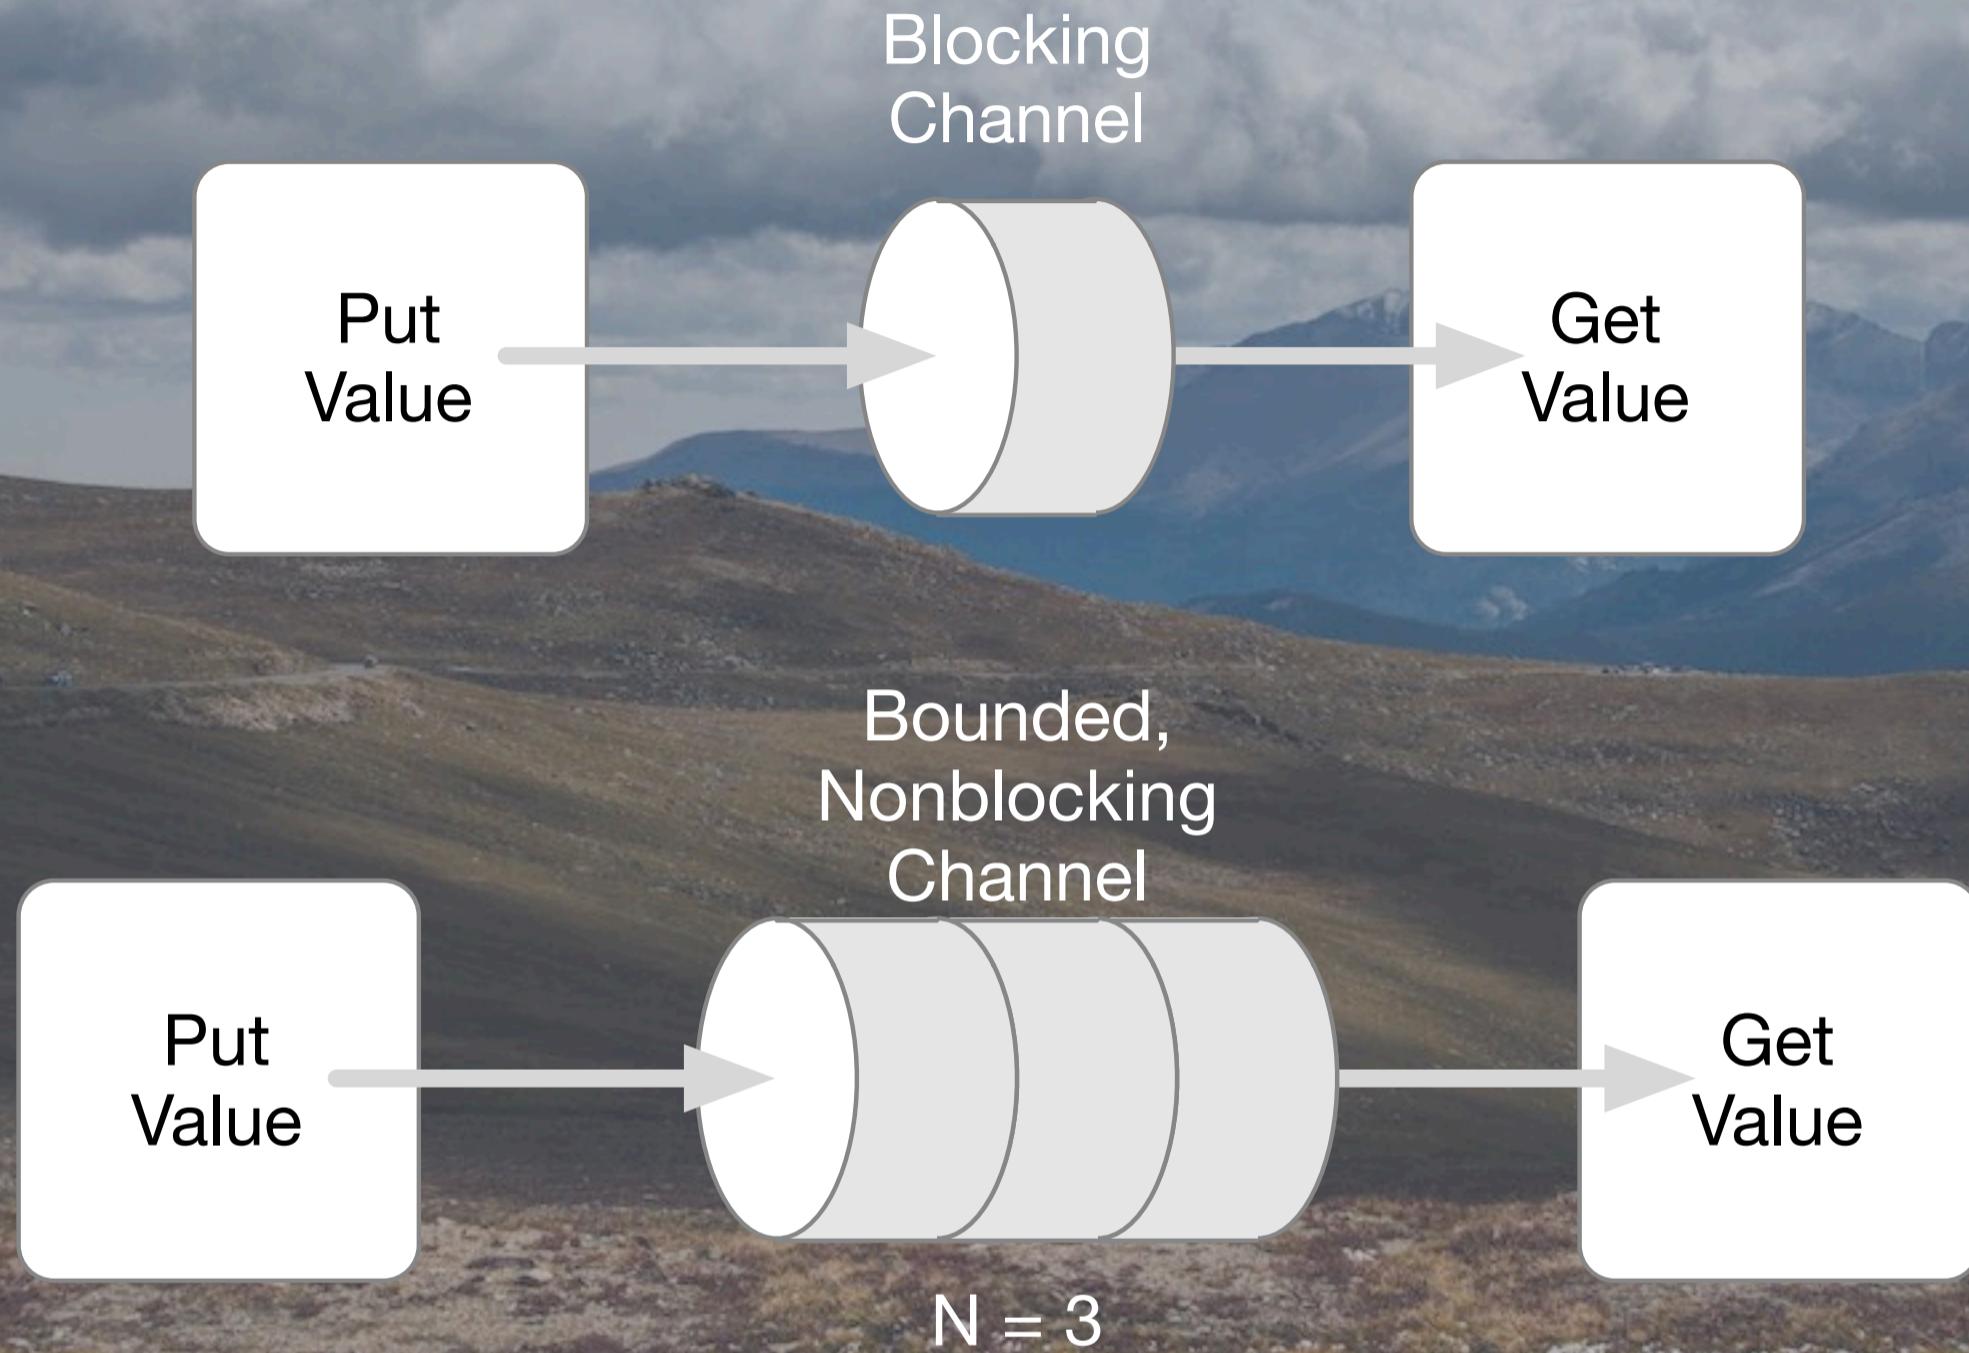
Tuesday, March 17, 15

Communicating Sequential Processes – The first mathematical model of distributed computing. It has evolved somewhat and it's still popular in Clojure and Go, for example.



Decouple sender and receiver
via a channel.
Can be sync. or async.
Not typically distributed...

CSP



Futures



87

Tuesday, March 17, 15

Fill more or less the same niche as CSP...

Futures

Define behaviors to run async.
Apply map, flatMap, ...
Or define success and
failure callbacks.

88

Tuesday, March 17, 15

Like using the UNIX shell to fire a process in the background, but you either define callbacks to handle the success or failure (a more procedure-oriented approach) or use functional operations like map, etc. to process the results on success.

Actor Systems

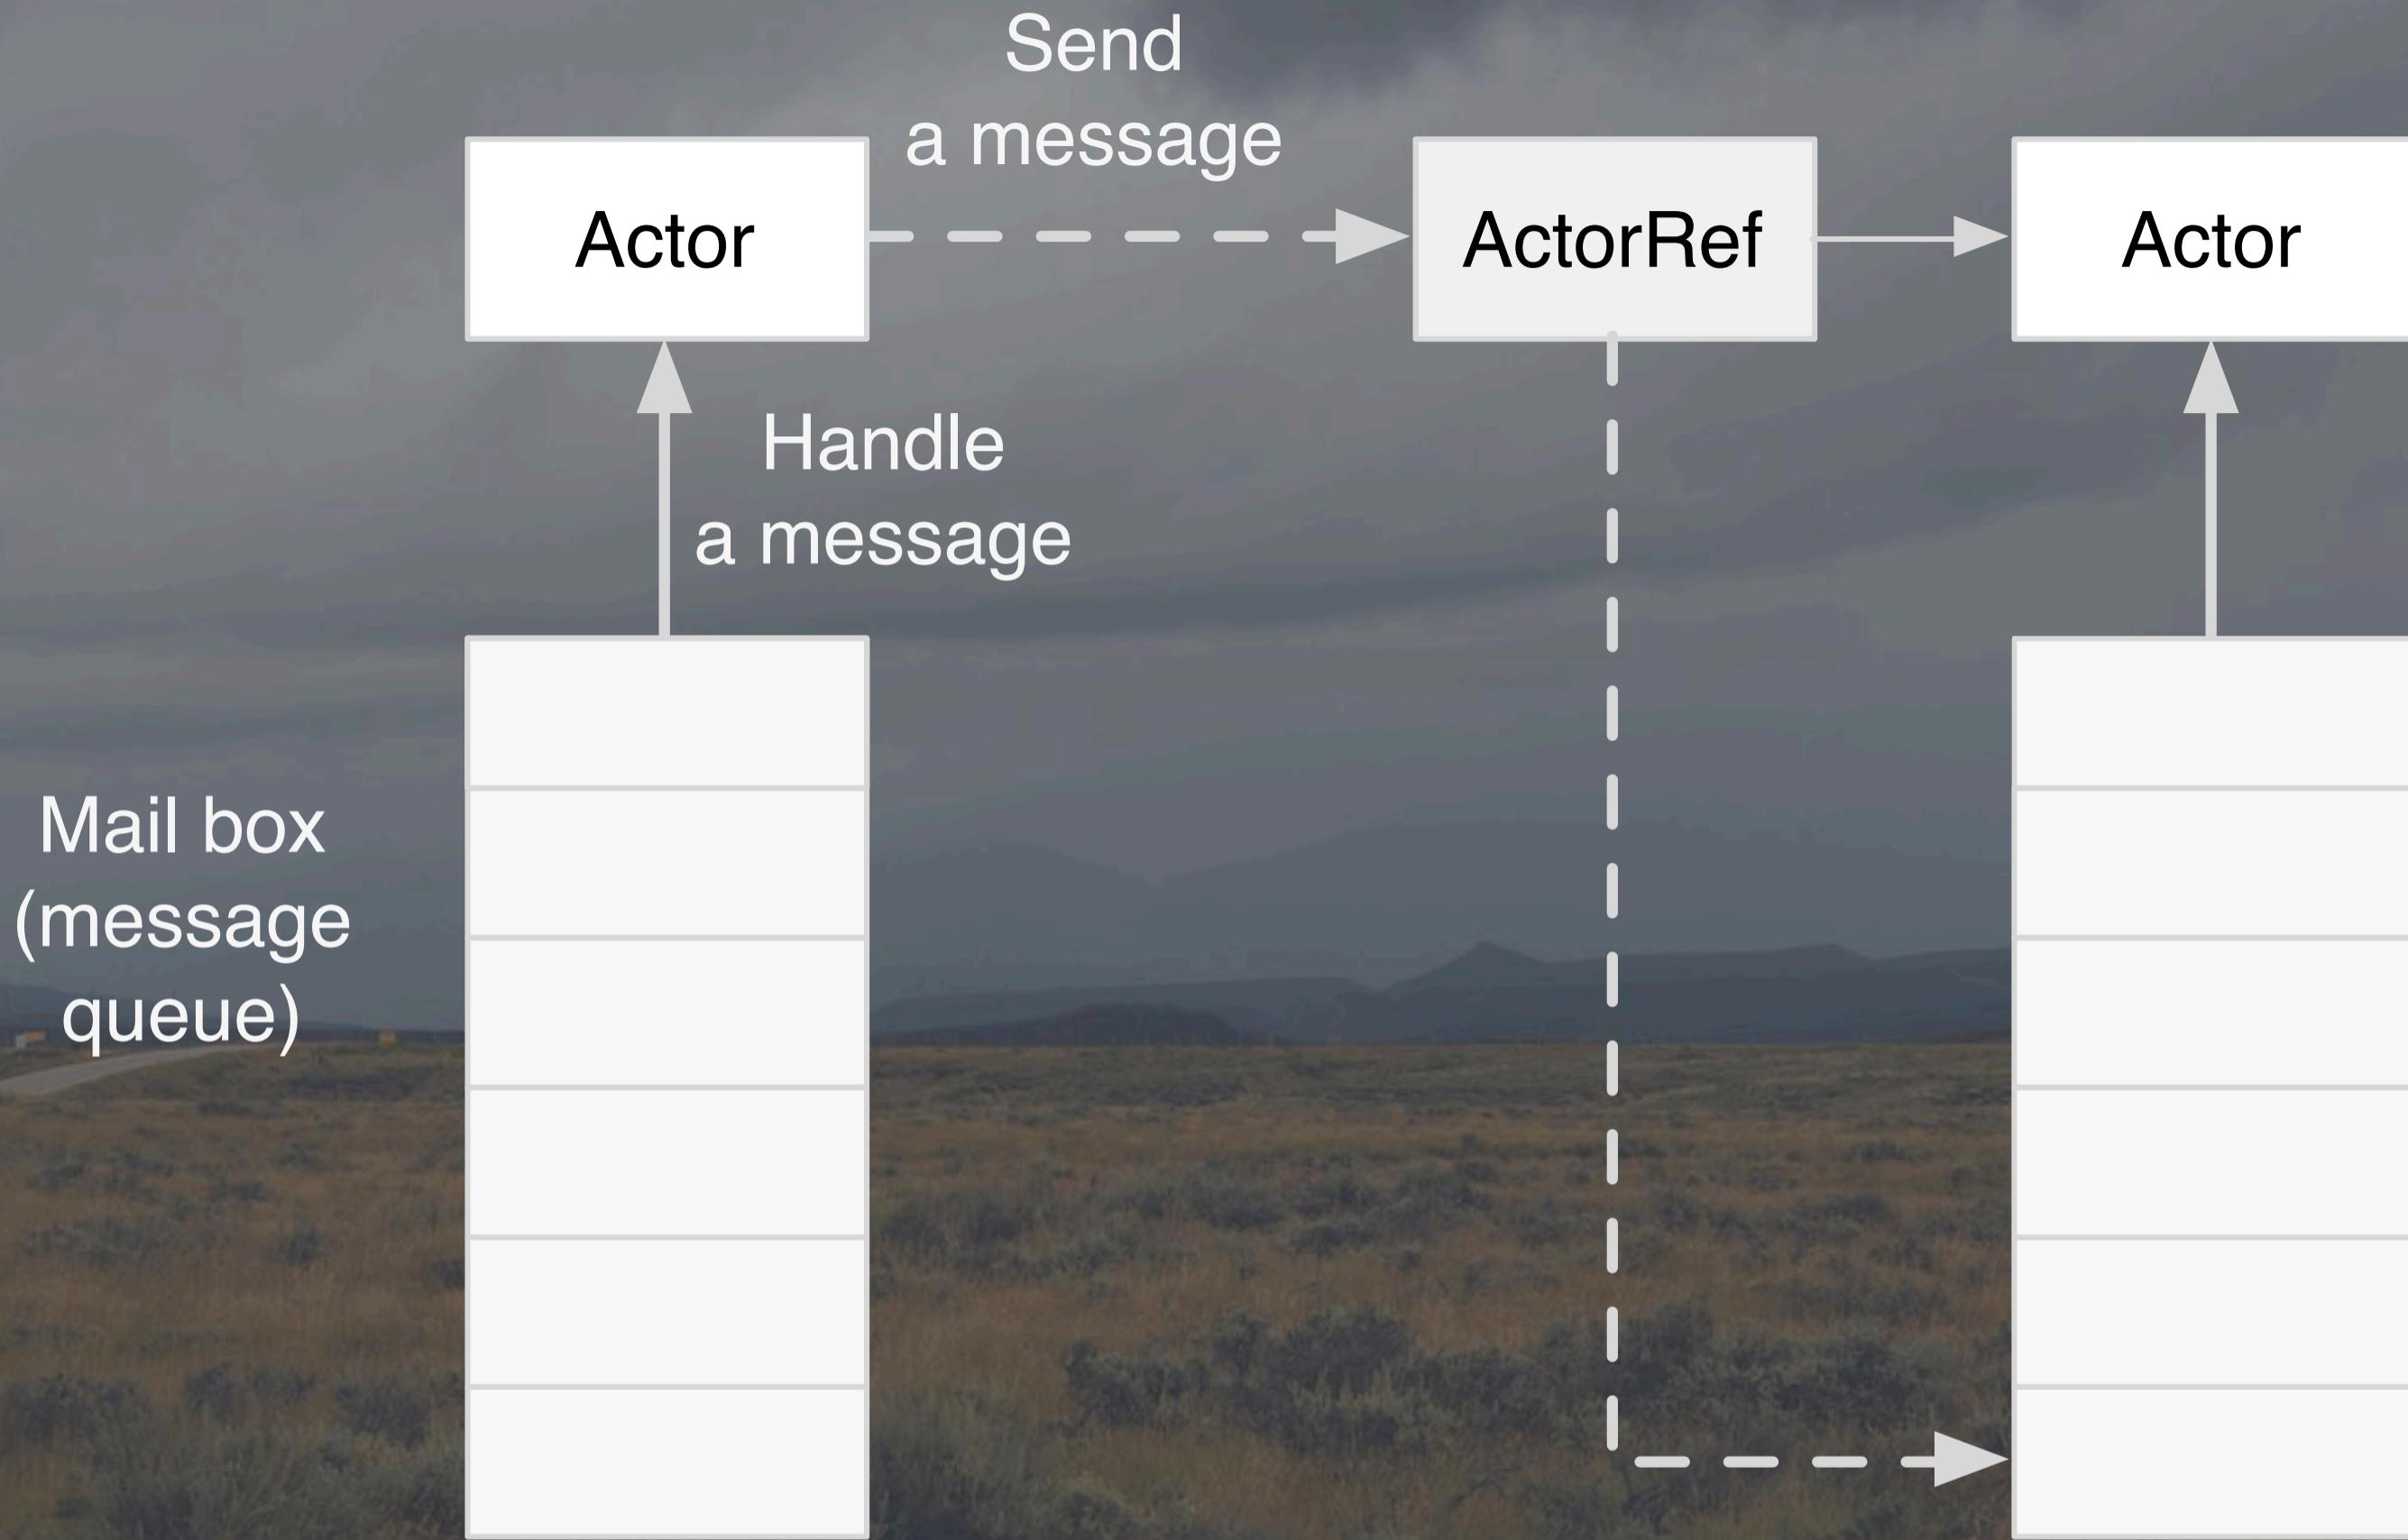
Let it Crash!

89

Tuesday, March 17, 15

We briefly visited this before.

Rather than attempt to recover from errors inside the domain logic (e.g., elaborate exception handling), allow services to fail, but with failure detection and reconstruction of those services, plus failover to other replicas.



90

Tuesday, March 17, 15

Actors are similar to objects in Smalltalk and similar systems; autonomous agents with defined boundaries that communicate through message passing. Actors, though process each message in a threadsafe way, so they are great for concurrency. (This diagram illustrates the Akka implementation – <http://akka.io>)

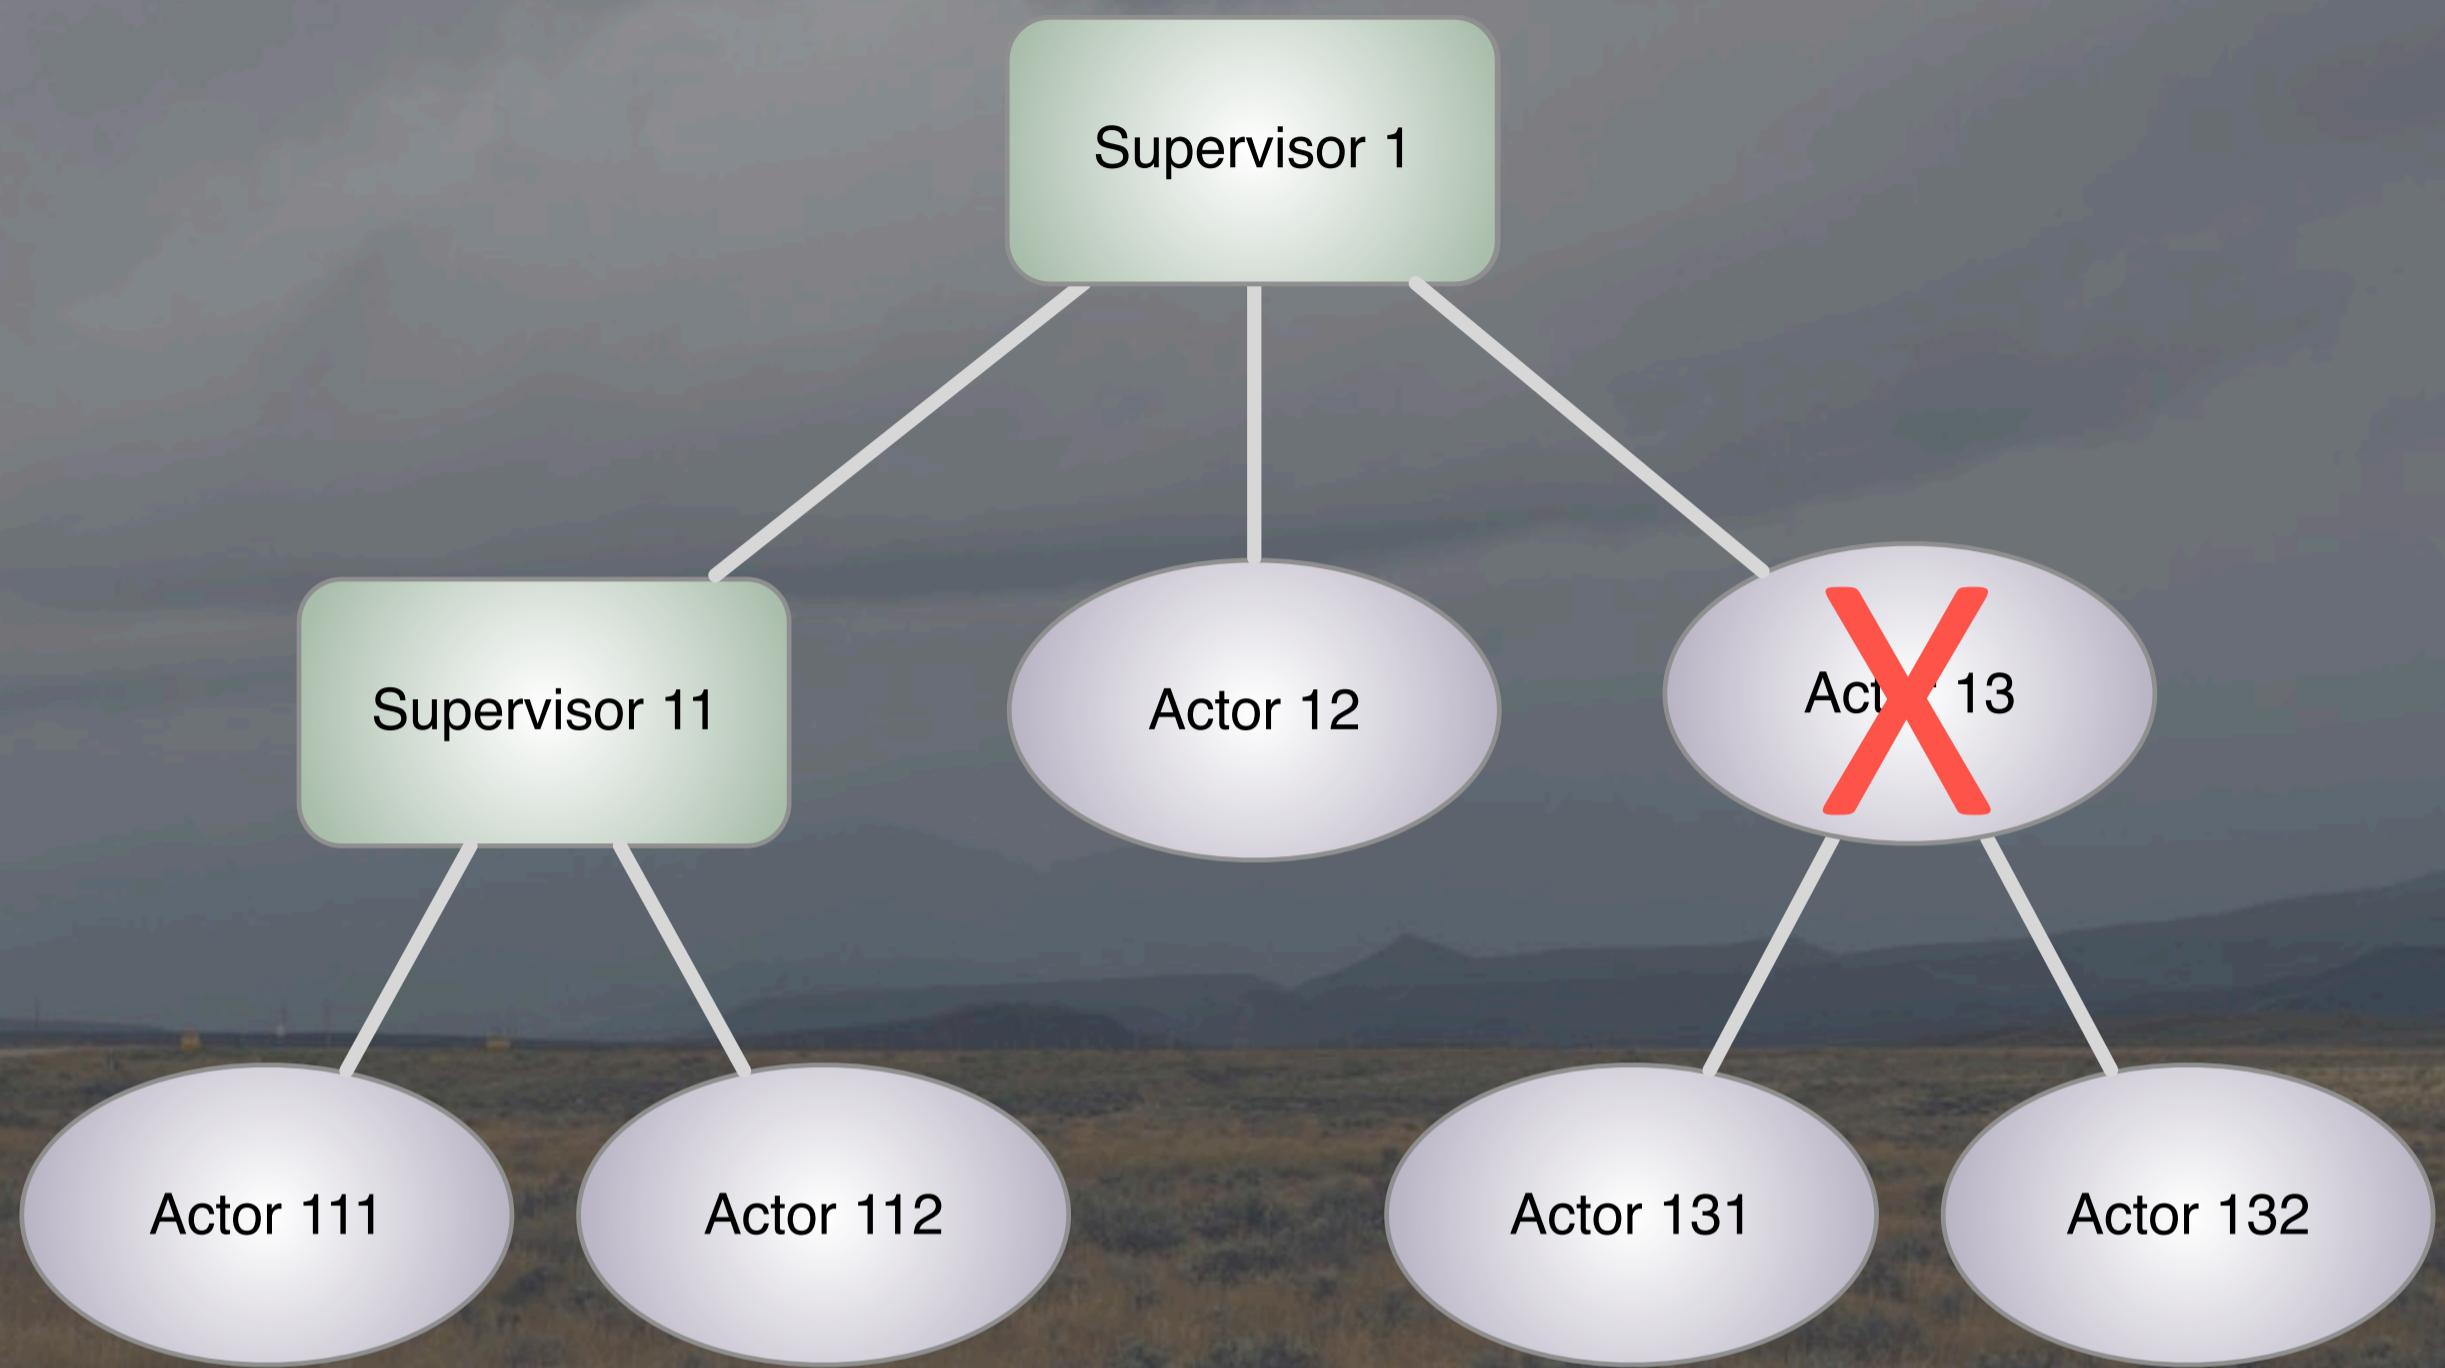
Actor Systems

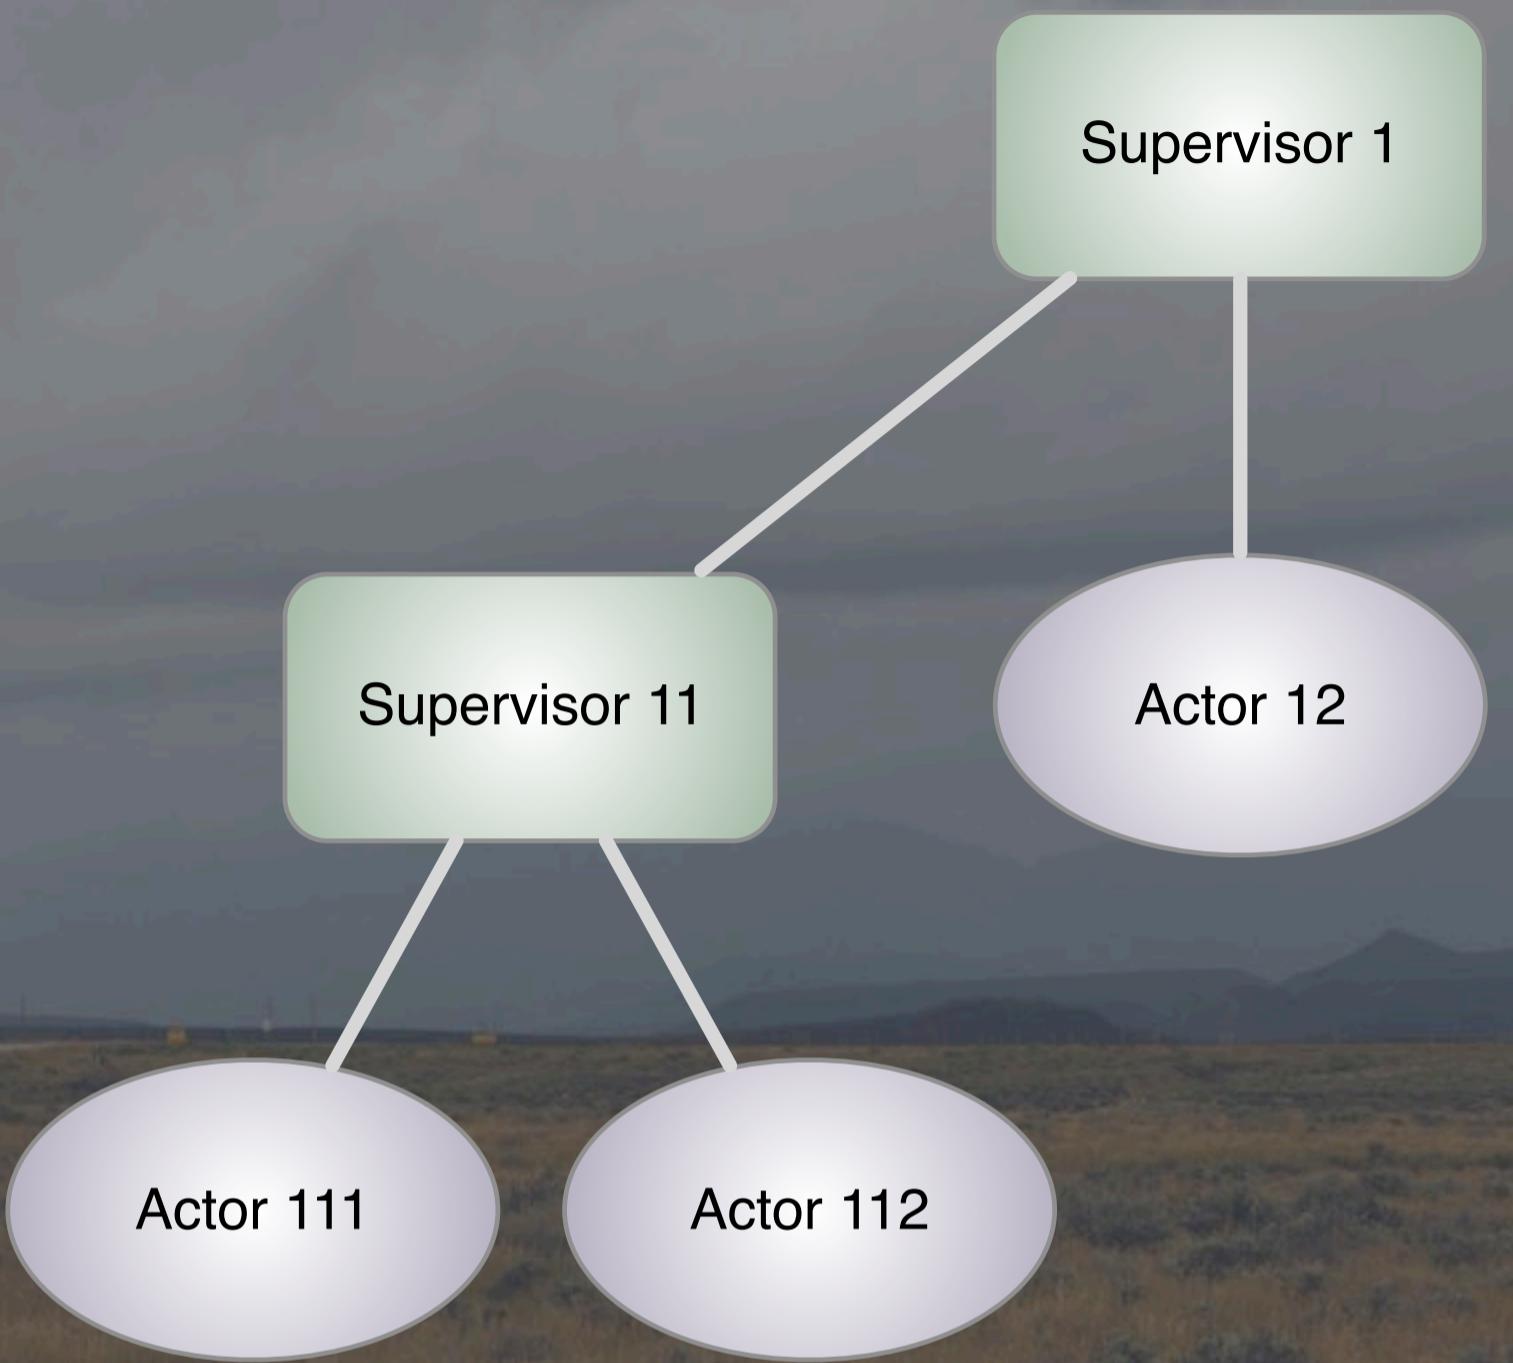
Pioneered by Hewitt, et al. 1973.
Made popular by Erlang,
which introduced Supervision.
Natural distribution model.

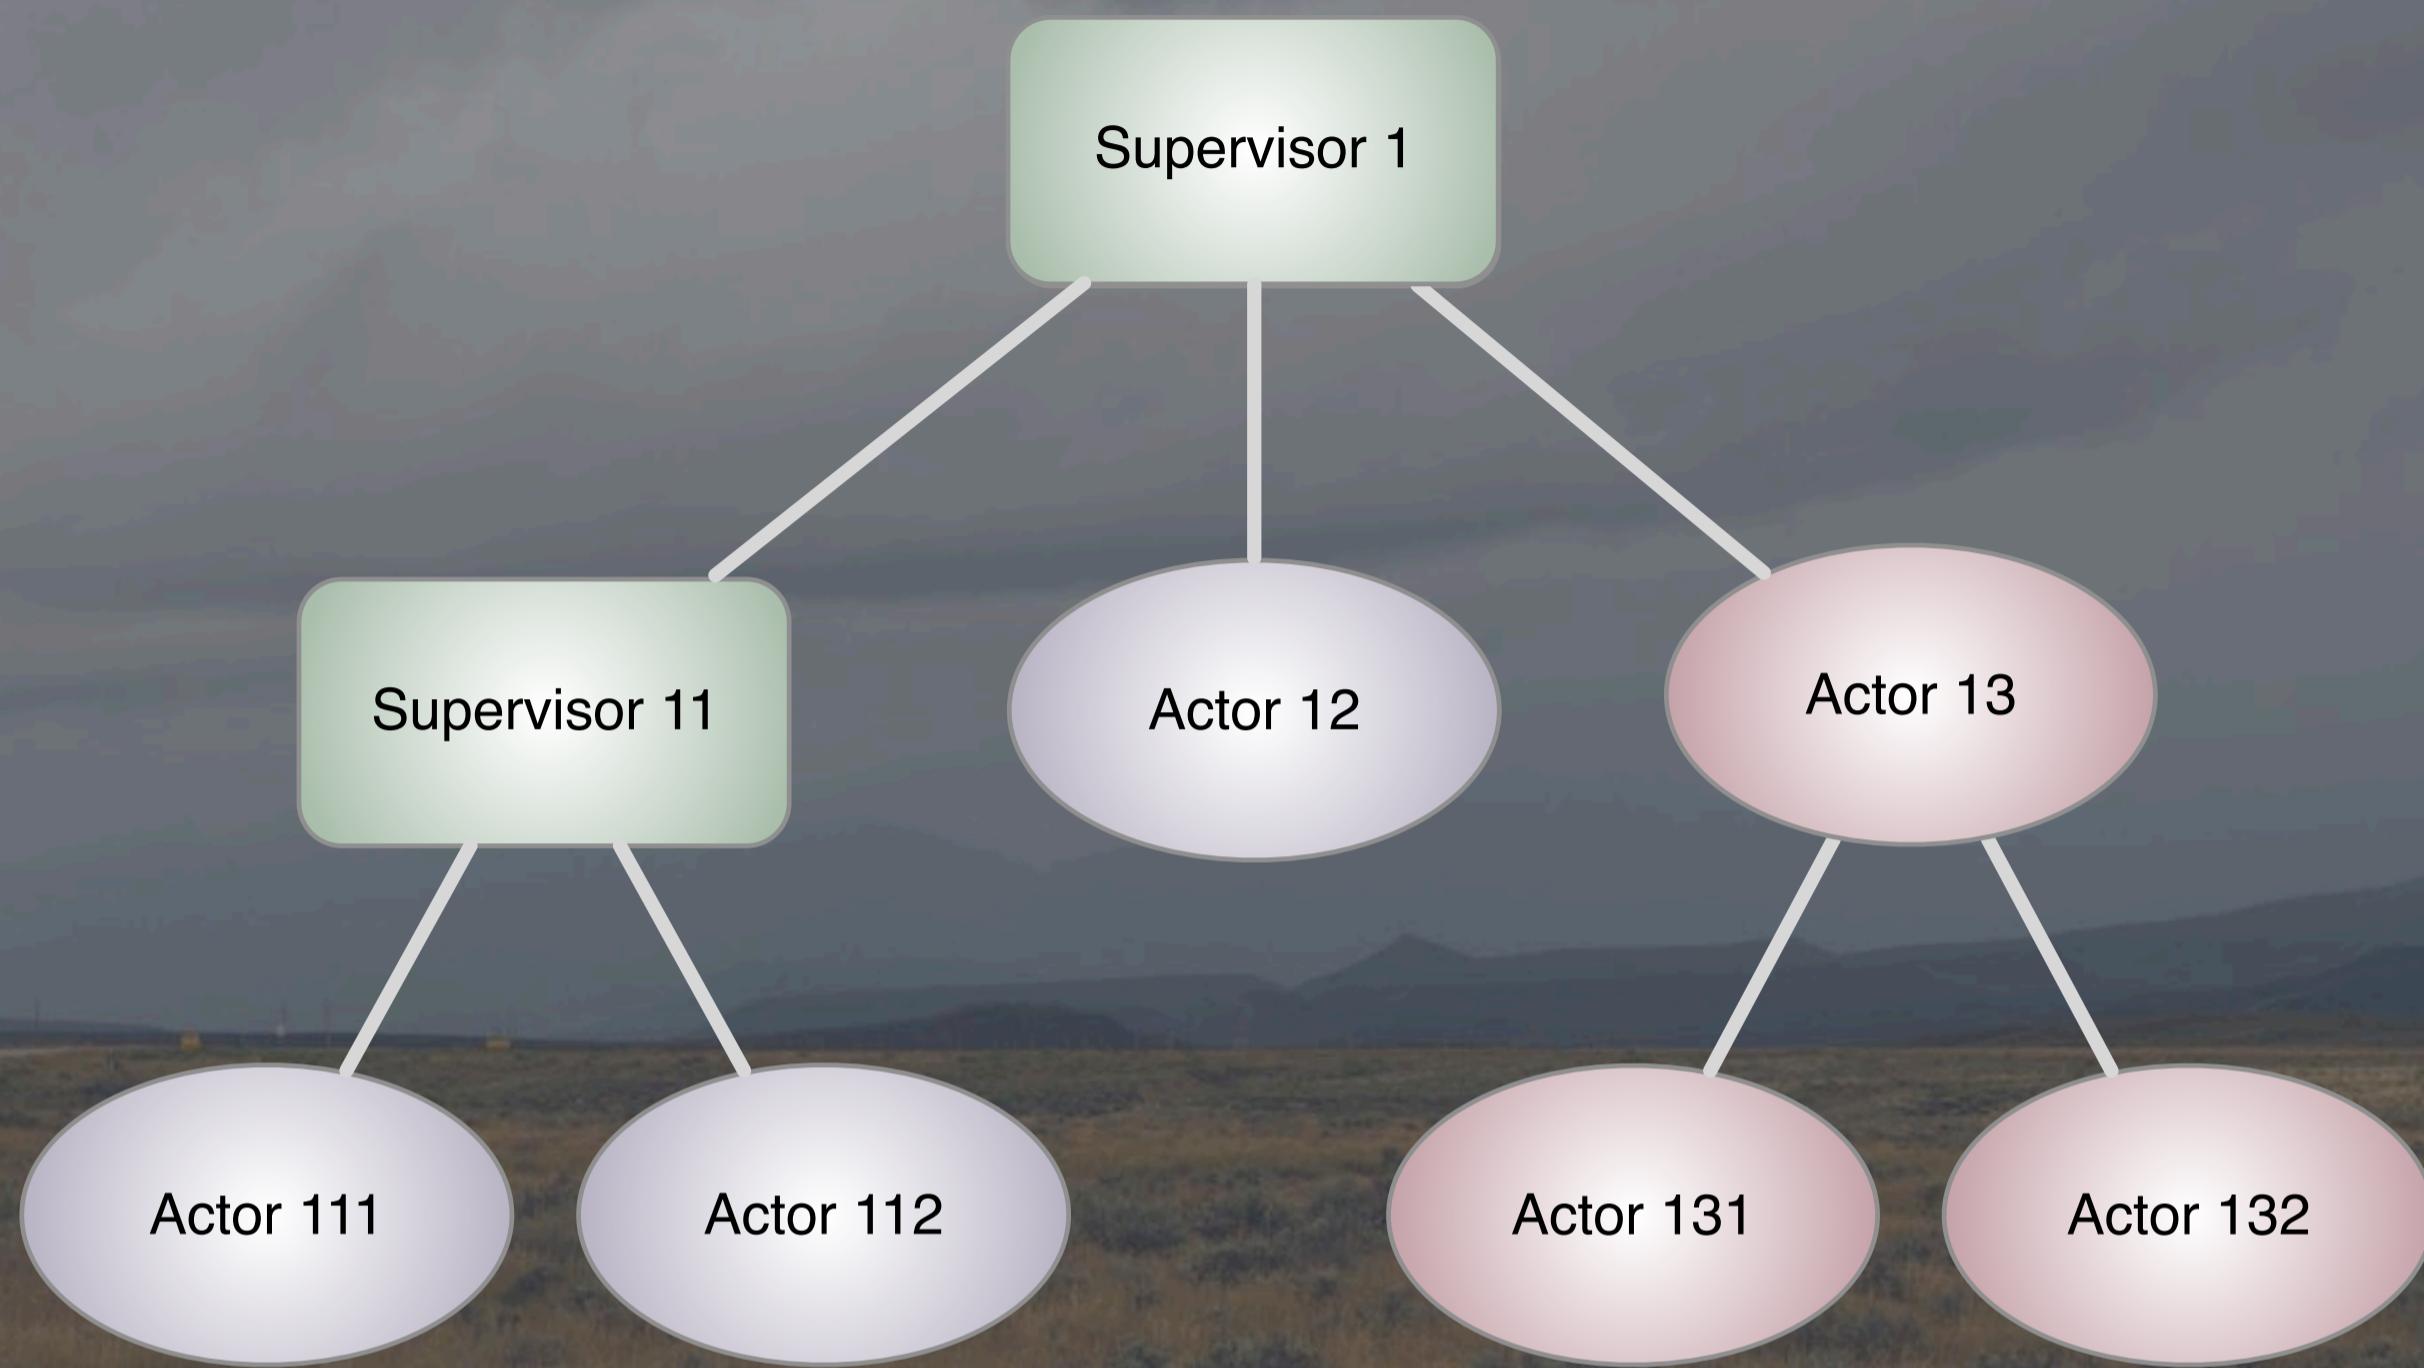
91

Tuesday, March 17, 15

Erlang is a simple language with actor semantics baked in. It has been used to create extremely reliable telecom switches, databases (e.g., Riak), and other services (e.g., GitHub).







A landscape photograph showing a dry, grassy field in the foreground, leading to a range of mountains in the background under a clear sky.

The most
sophisticated error recovery
in reactive systems.

95



Clean separation
of normal processing
from recovery.

A landscape photograph showing a dry, grassy field in the foreground, leading to a range of mountains in the background under a clear sky.

State mutation “firewalls”.
Supports location transparency.

Reactive Extensions (Rx)

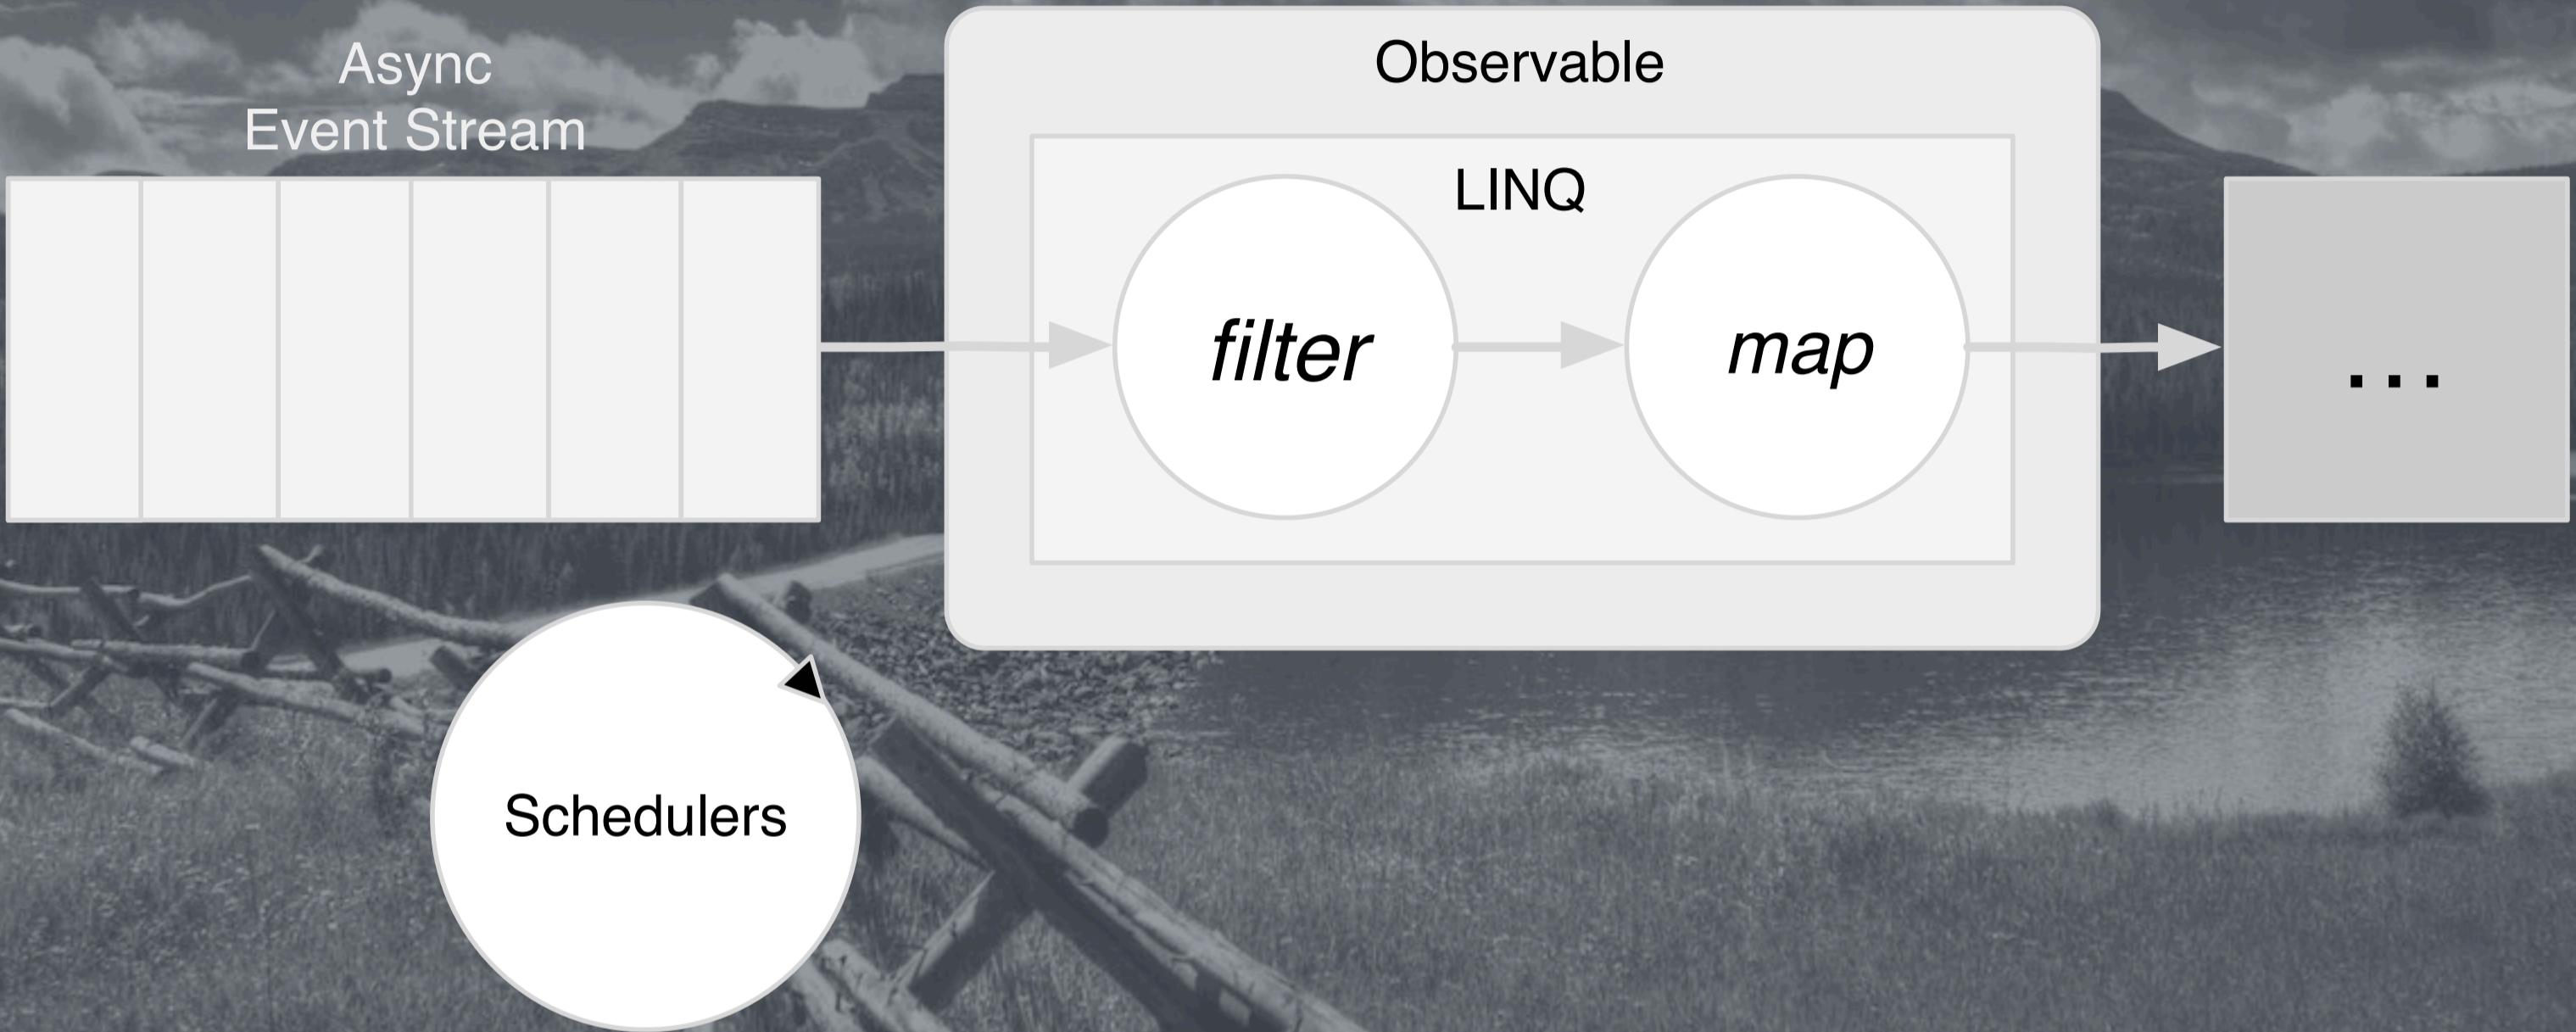


98

Tuesday, March 17, 15

Pioneered by Erik Meijer for .NET. Now ported to several languages, including RxJava (Netflix) and React (JS – Facebook).

Rx



Rx

Events pushed to system.
FP/SQL-like query
semantics simplify
handling logic.

100

Rx

Combines Iterator and
Observer into Observable.
Stream oriented.

101

Rx

Need to add your own
fault-tolerance model.

Reactive Streams

reactive-streams.org

103

Tuesday, March 17, 15

A standard with many implementations for streaming systems with truly “reactive” behavior.
photo: Bridge Creek, North Cascades National Park, Washington State (not Colorado ;)

Reactive Streams

Decouples processing.
Adds back pressure,
focuses on FP
transformations.

104

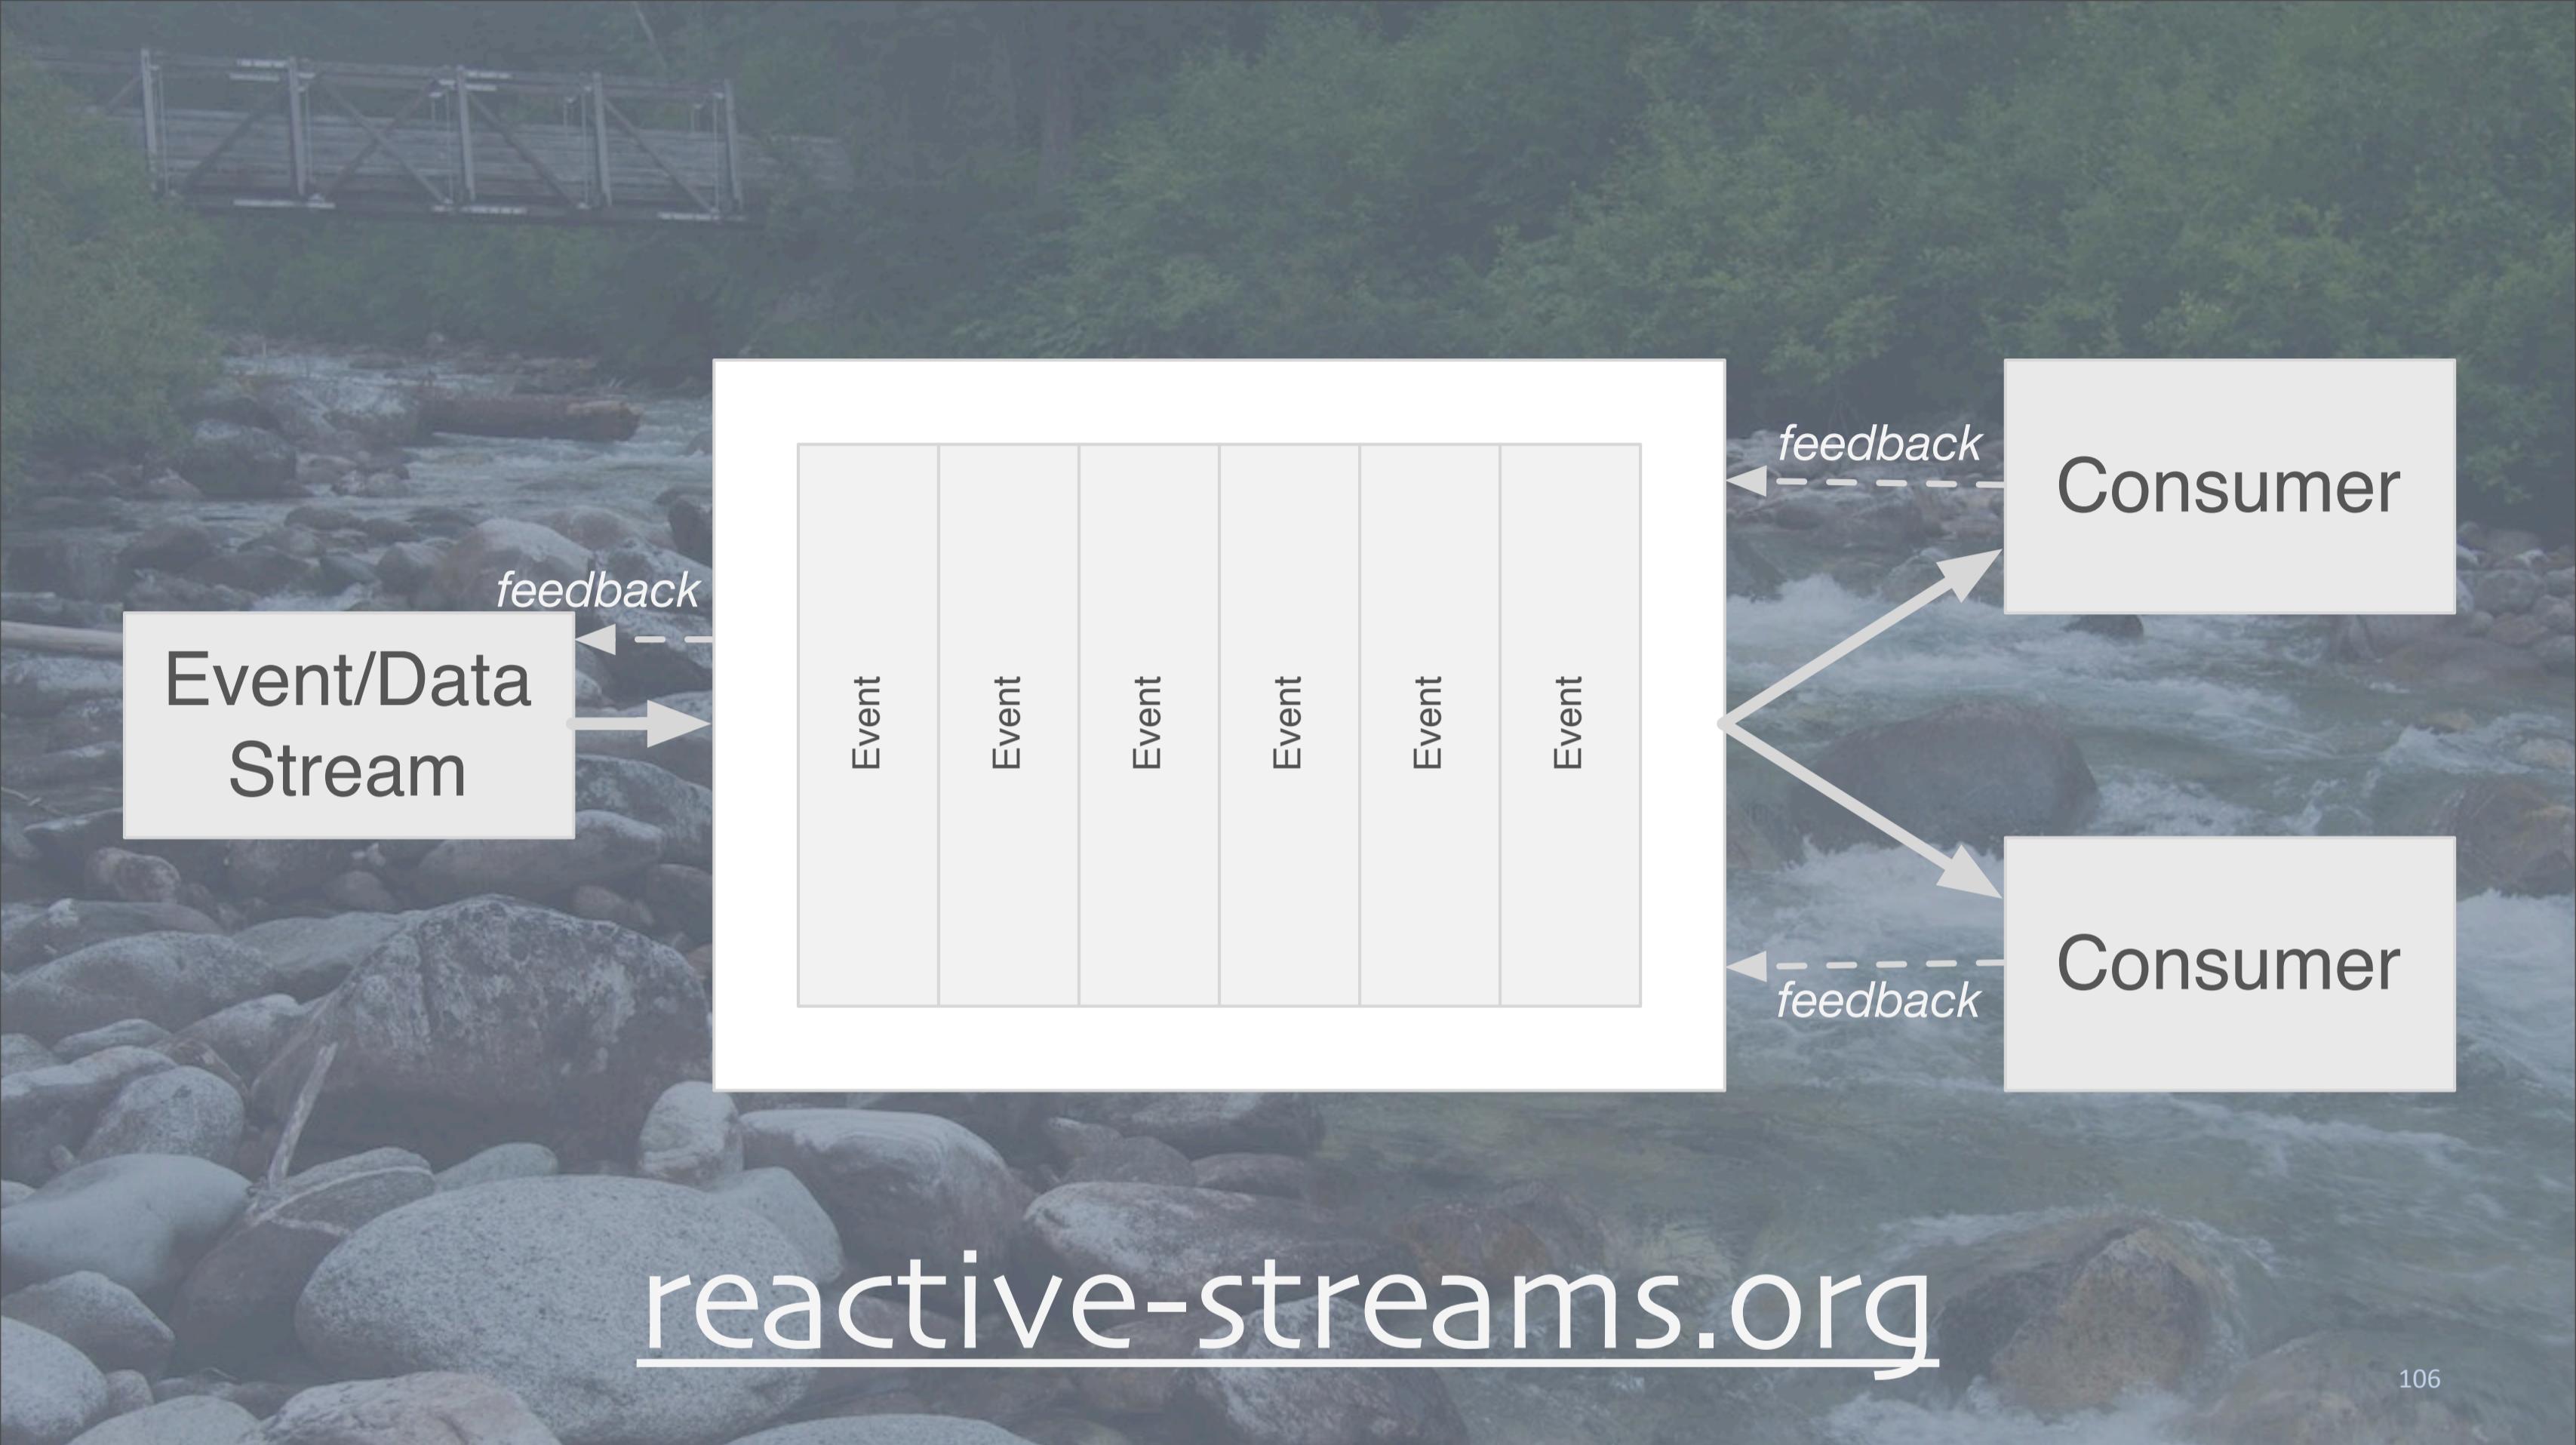
Tuesday, March 17, 15

Back pressure, where the producer and consumer negotiate the flow rate dynamically, is essential for avoiding many failure scenarios, like buffer overflows, ad hoc dropping of messages, etc.

Reactive Streams

Logical extension of Rx.
Focus on possibly
infinite streams of Data.

105



reactive-streams.org

106



Reactive Streams

Akka Streams: a higher-level abstraction implemented with Akka Actors.

107

Tuesday, March 17, 15

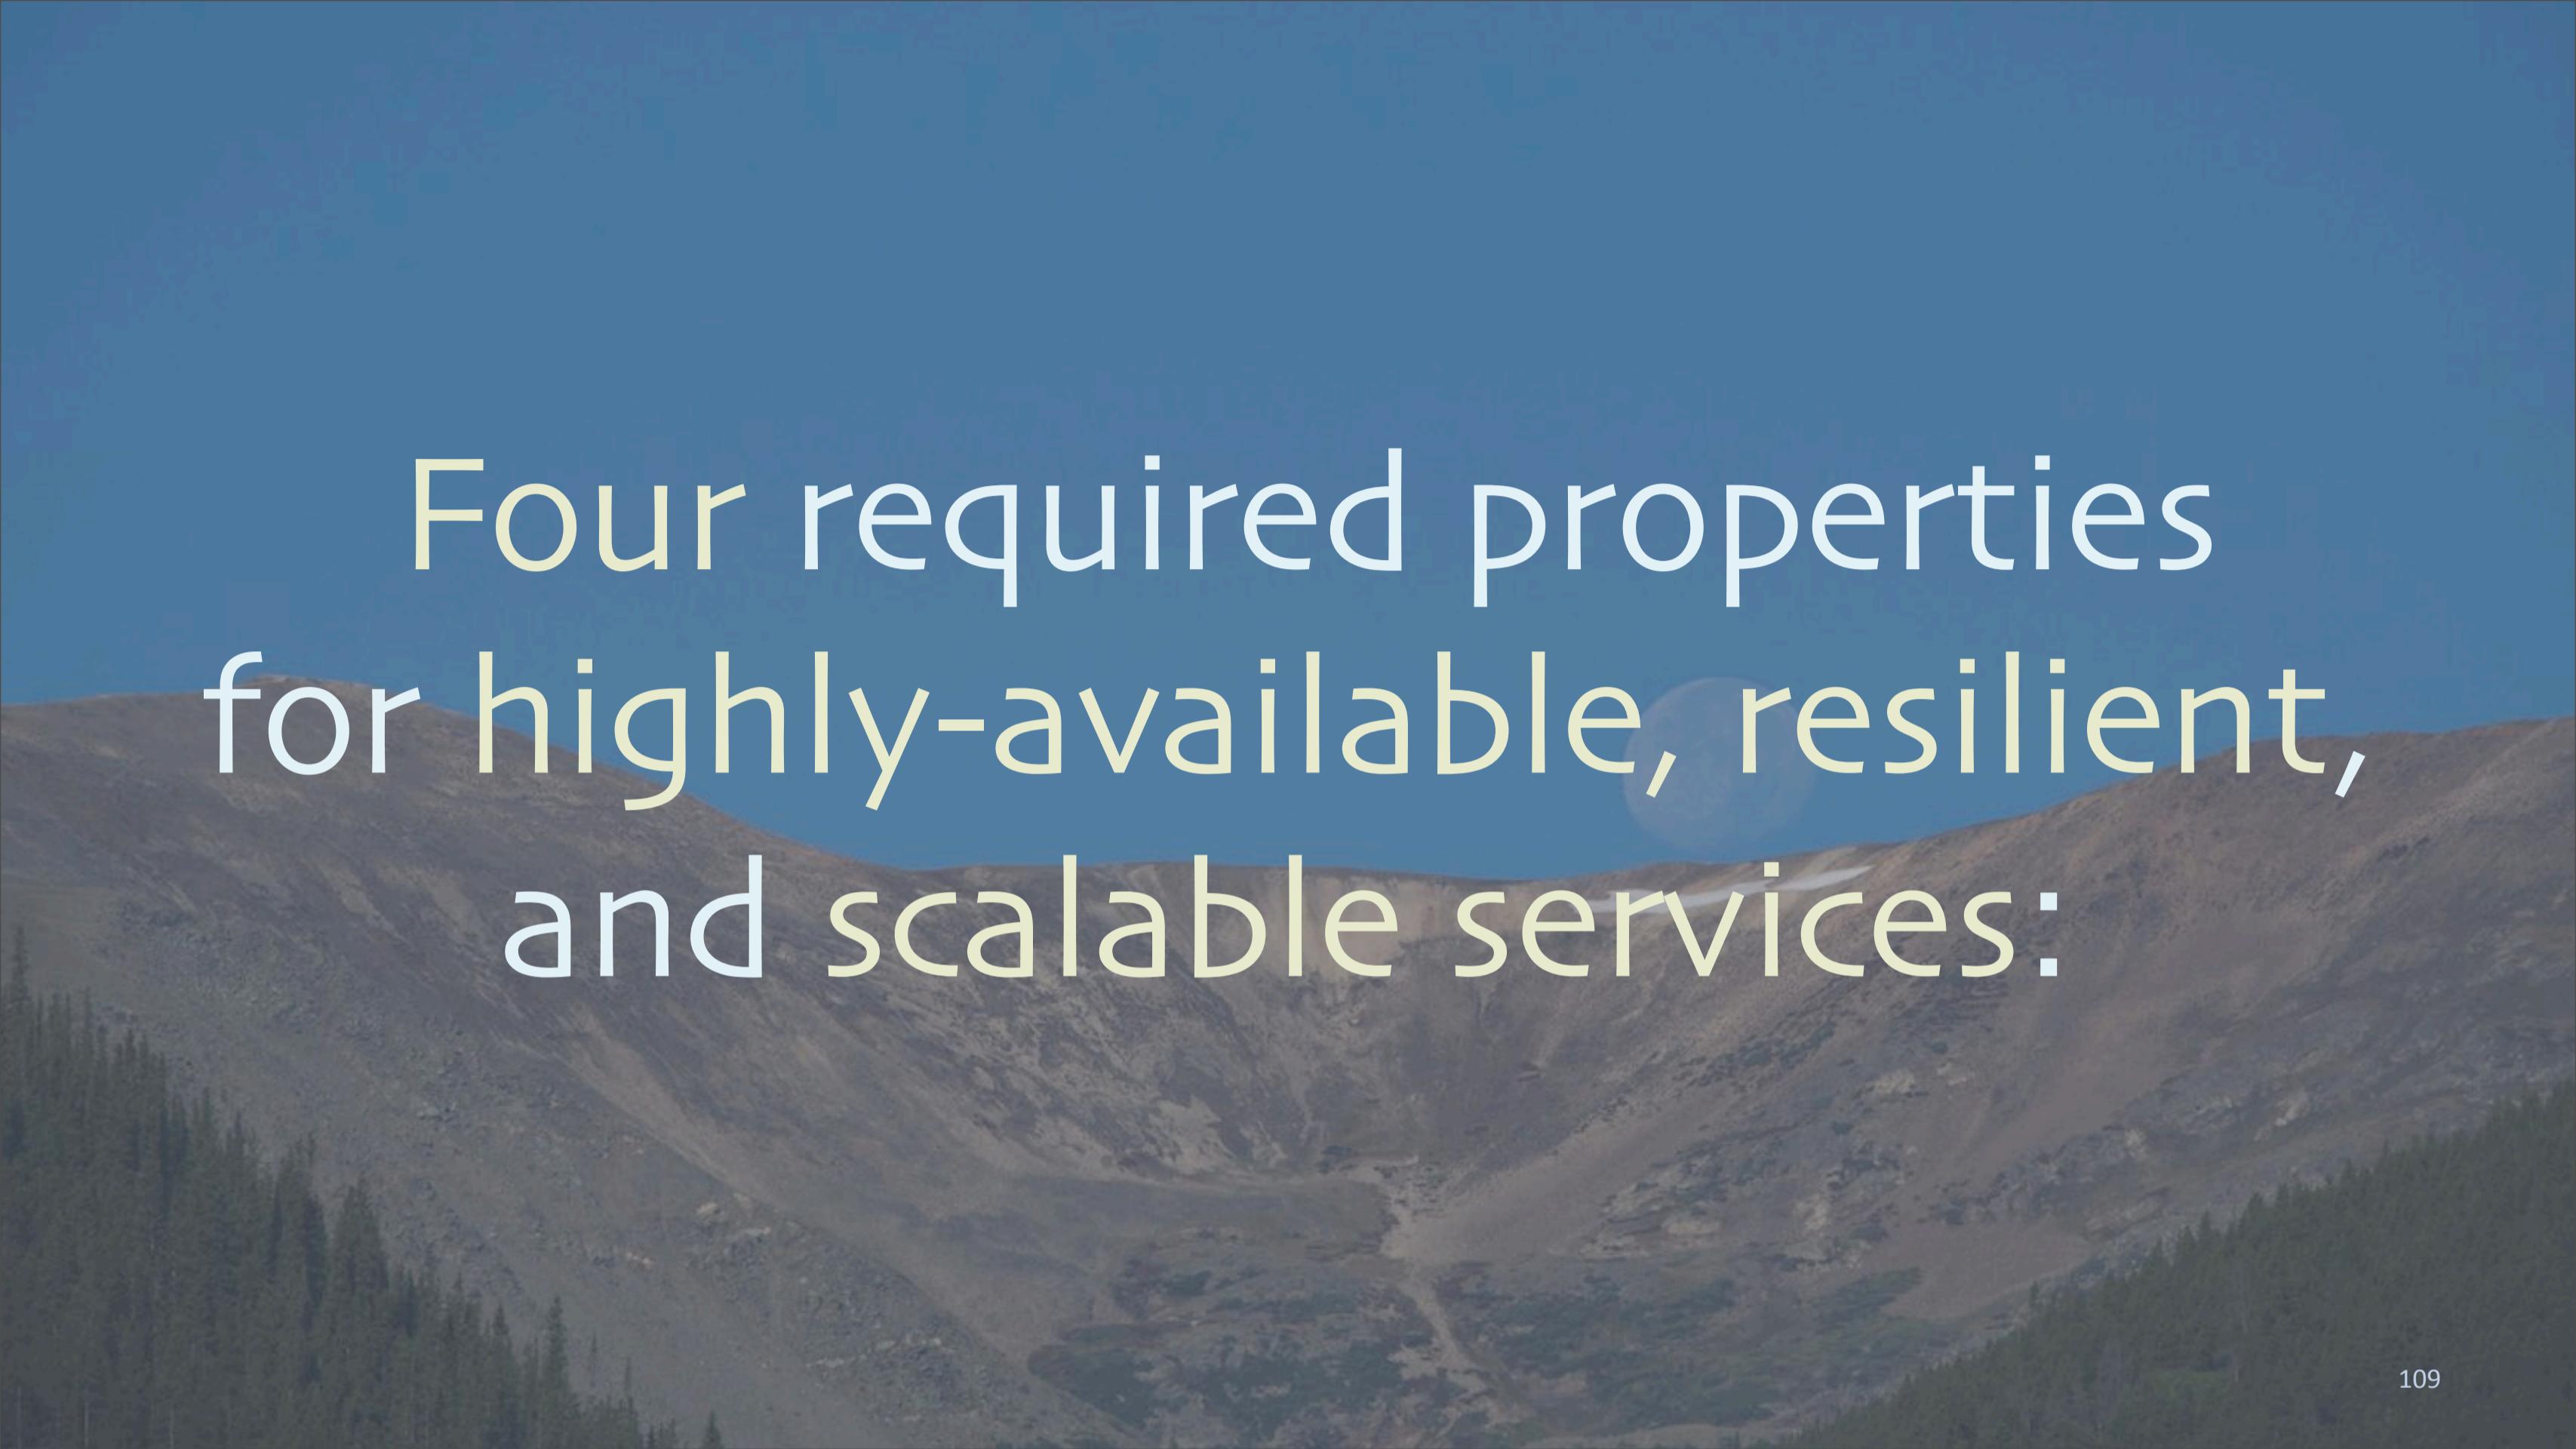
We're realizing that Actors are a low-level primitive and Actor systems can become unwieldy. Typesafe thinks that higher-level abstractions, like reactive streams, implemented on low-level concurrency systems, like Actors, will be the way to go for most future systems.

Recap

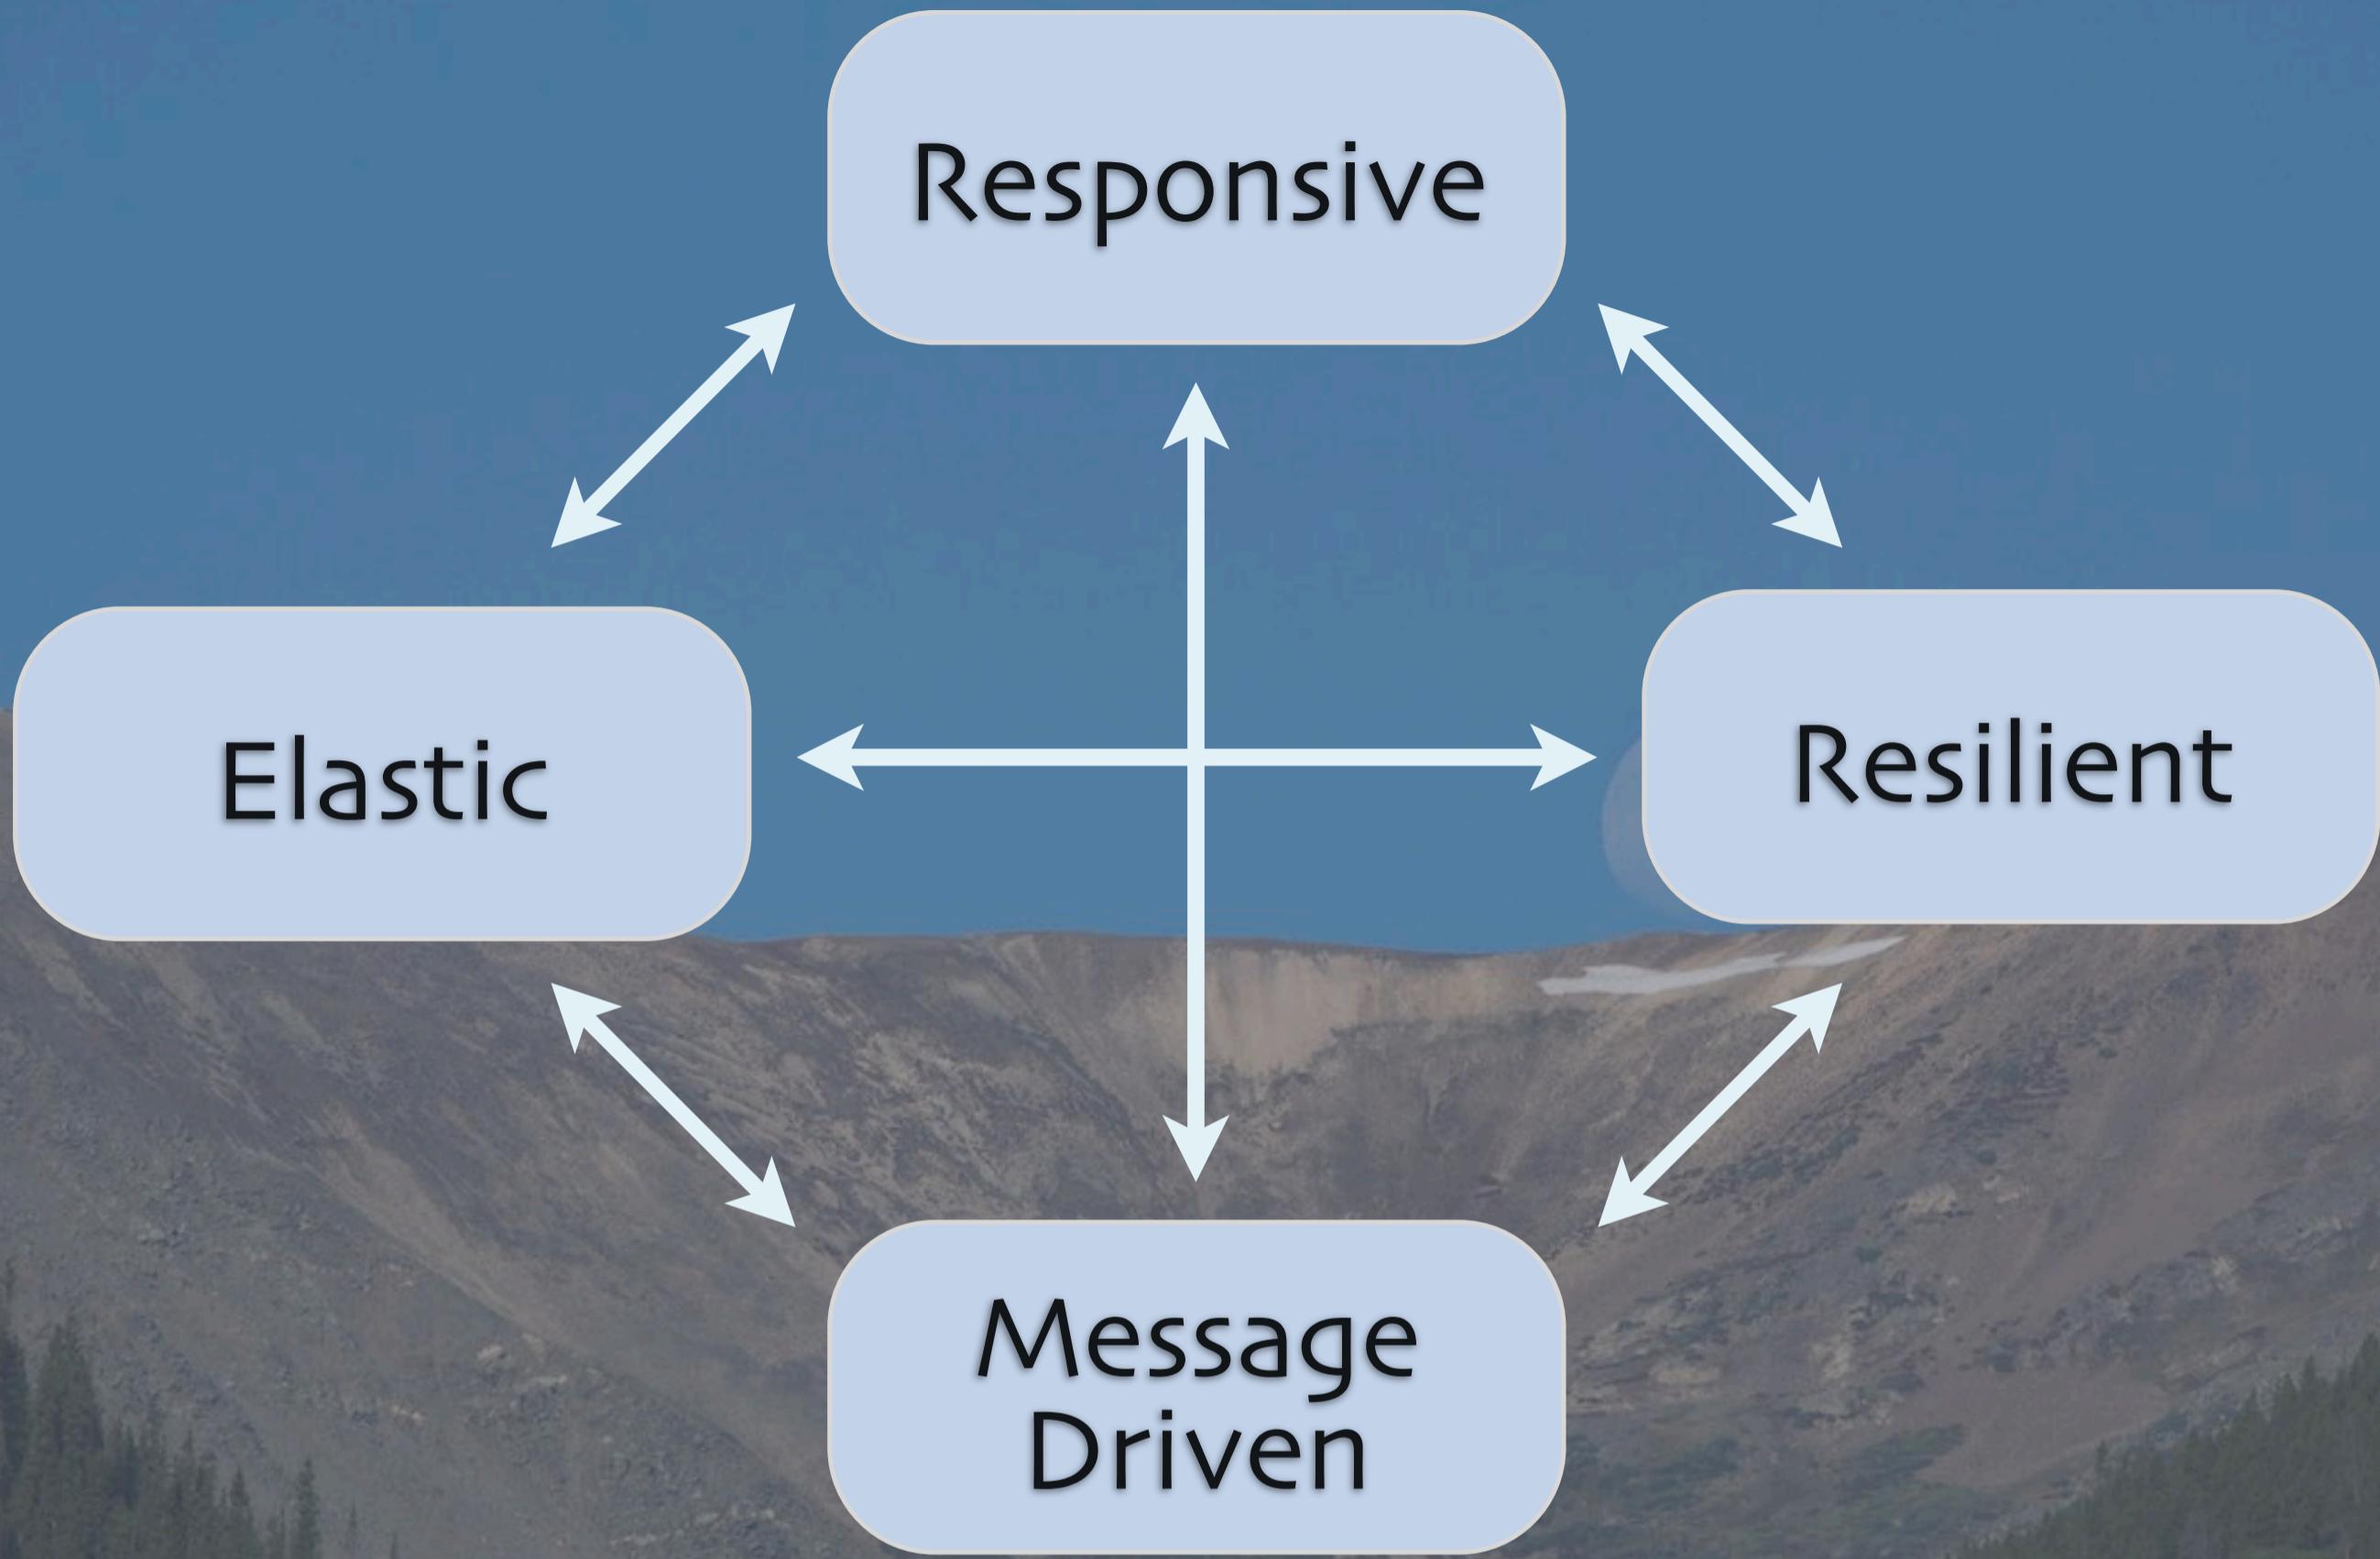


108

Tuesday, March 17, 15



Four required properties
for highly-available, resilient,
and scalable services:





Typesafe

Reactive

<http://typesafe.com/reactive-big-data>

dean.wampler@typesafe.com

©Typesafe 2014-2015, All Rights Reserved

Tuesday, March 17, 15

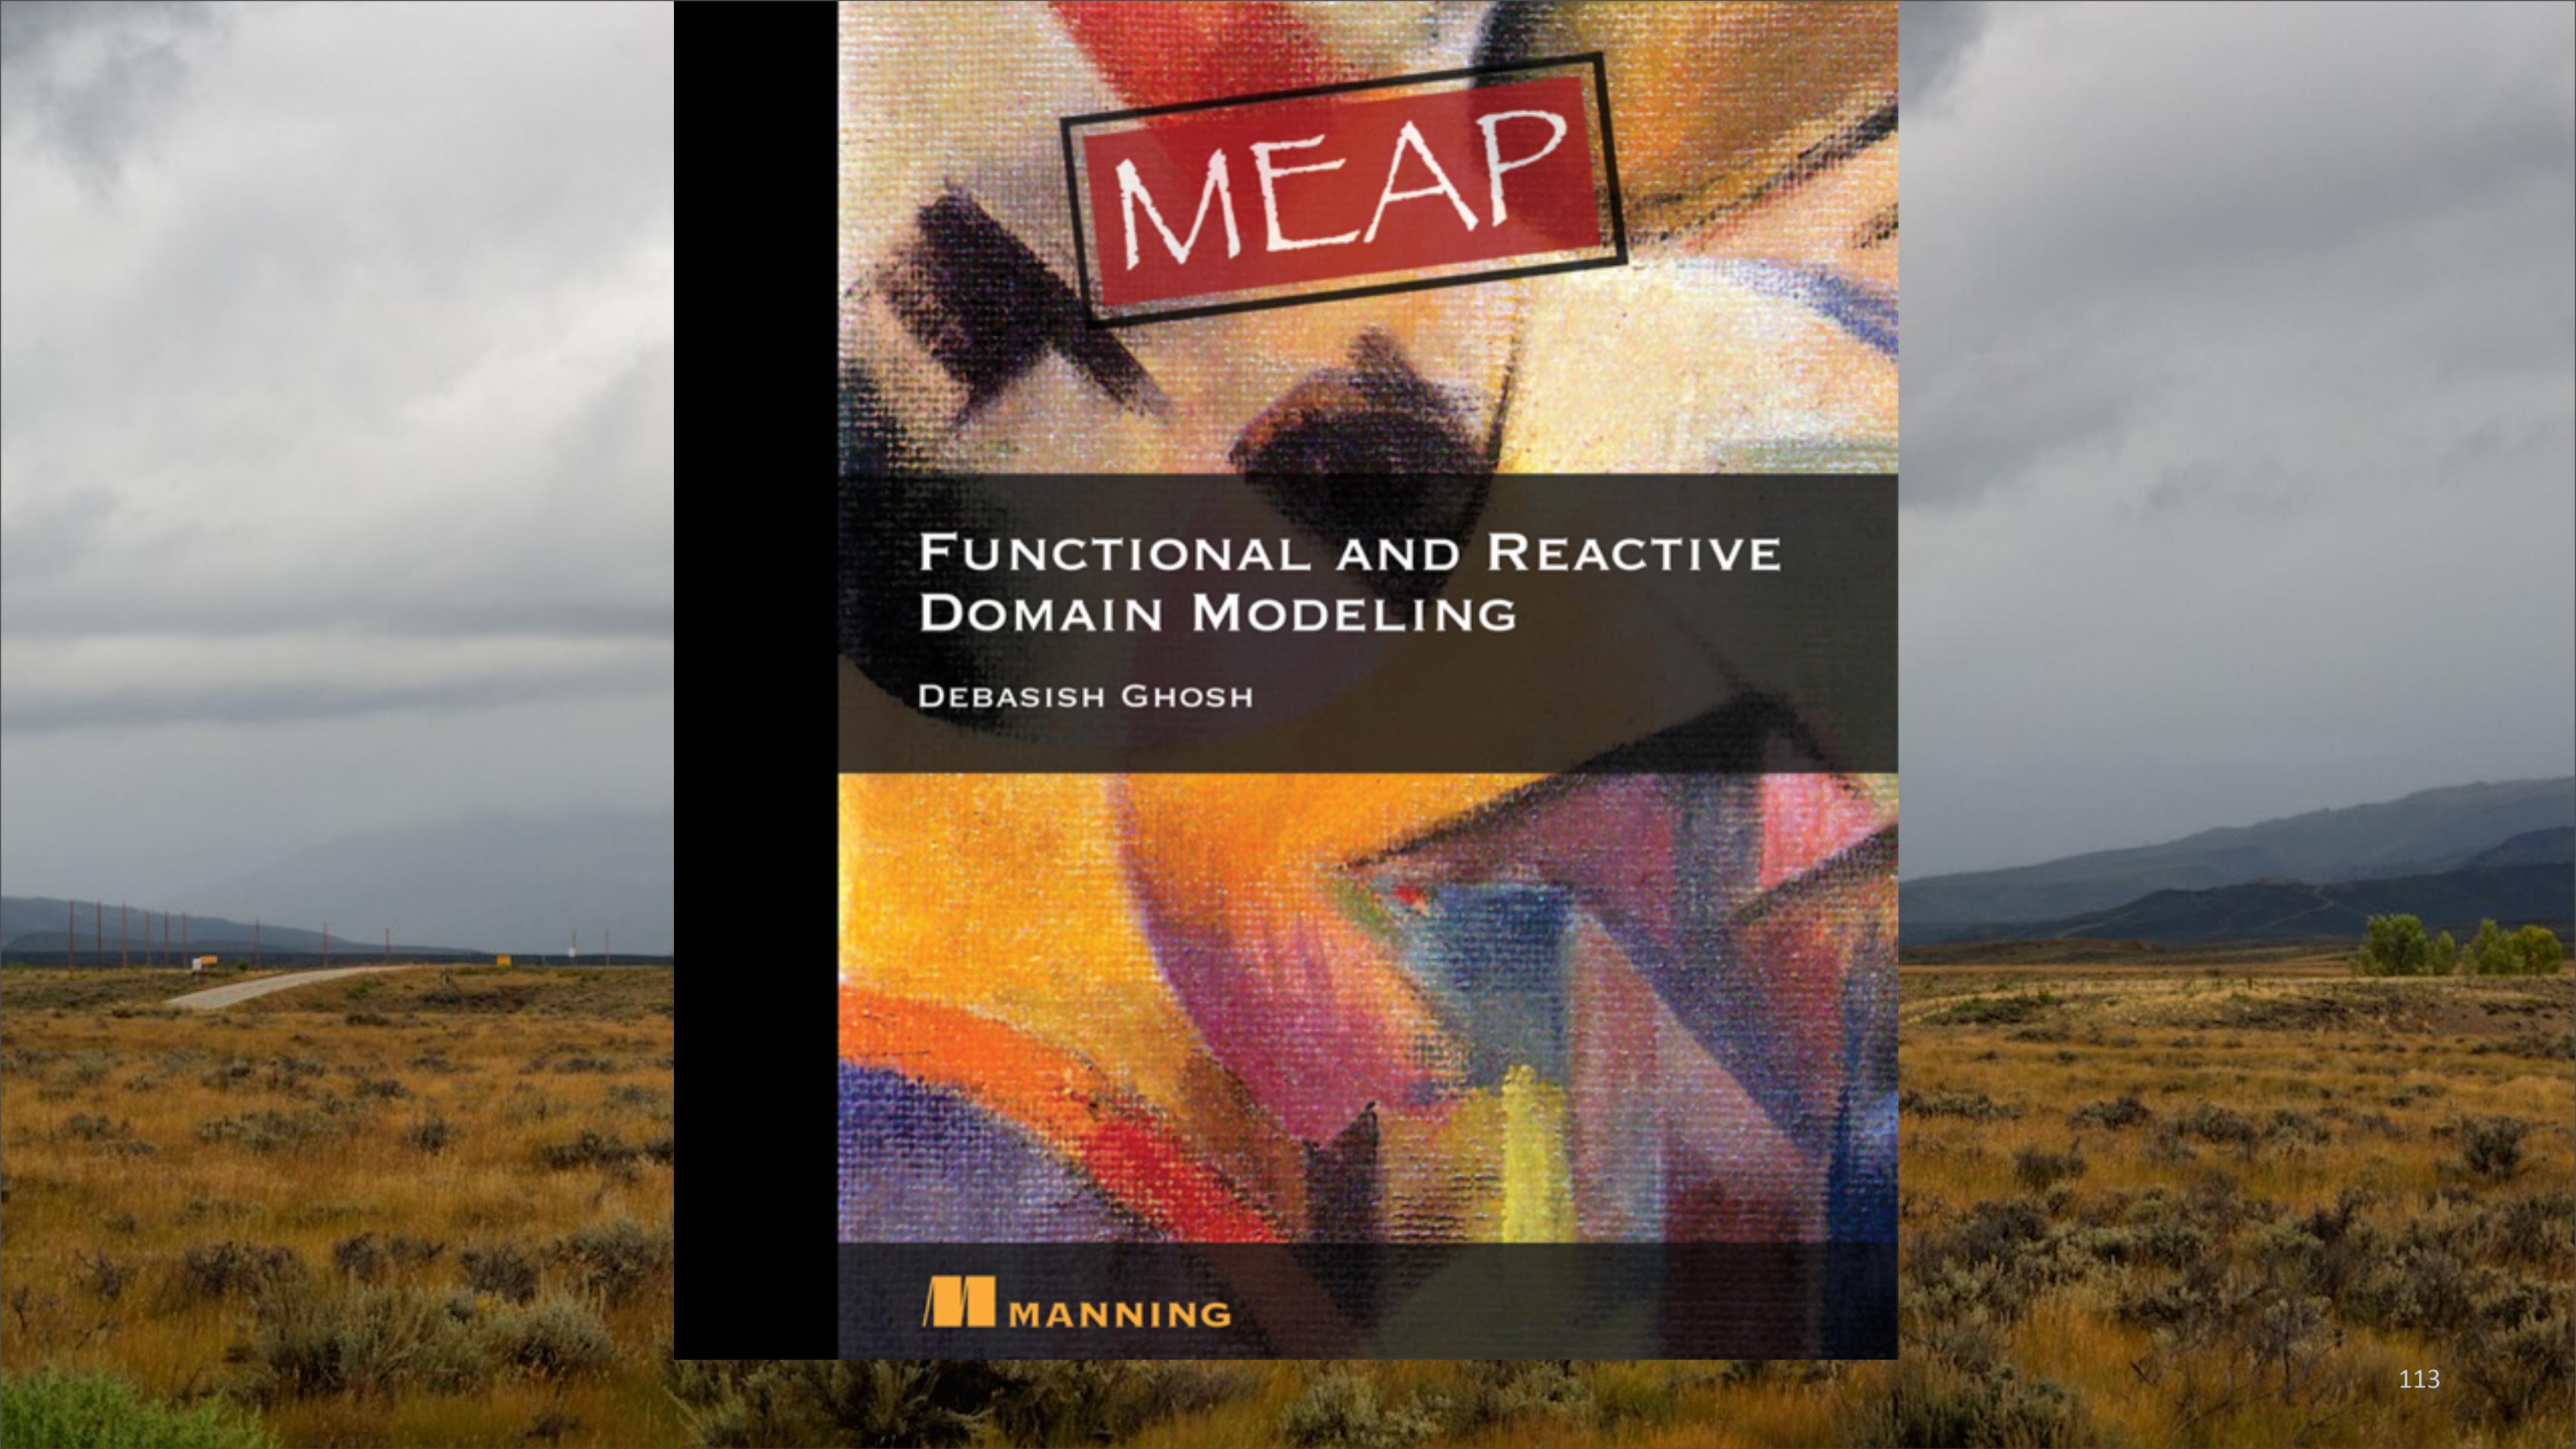


References

112

Tuesday, March 17, 15

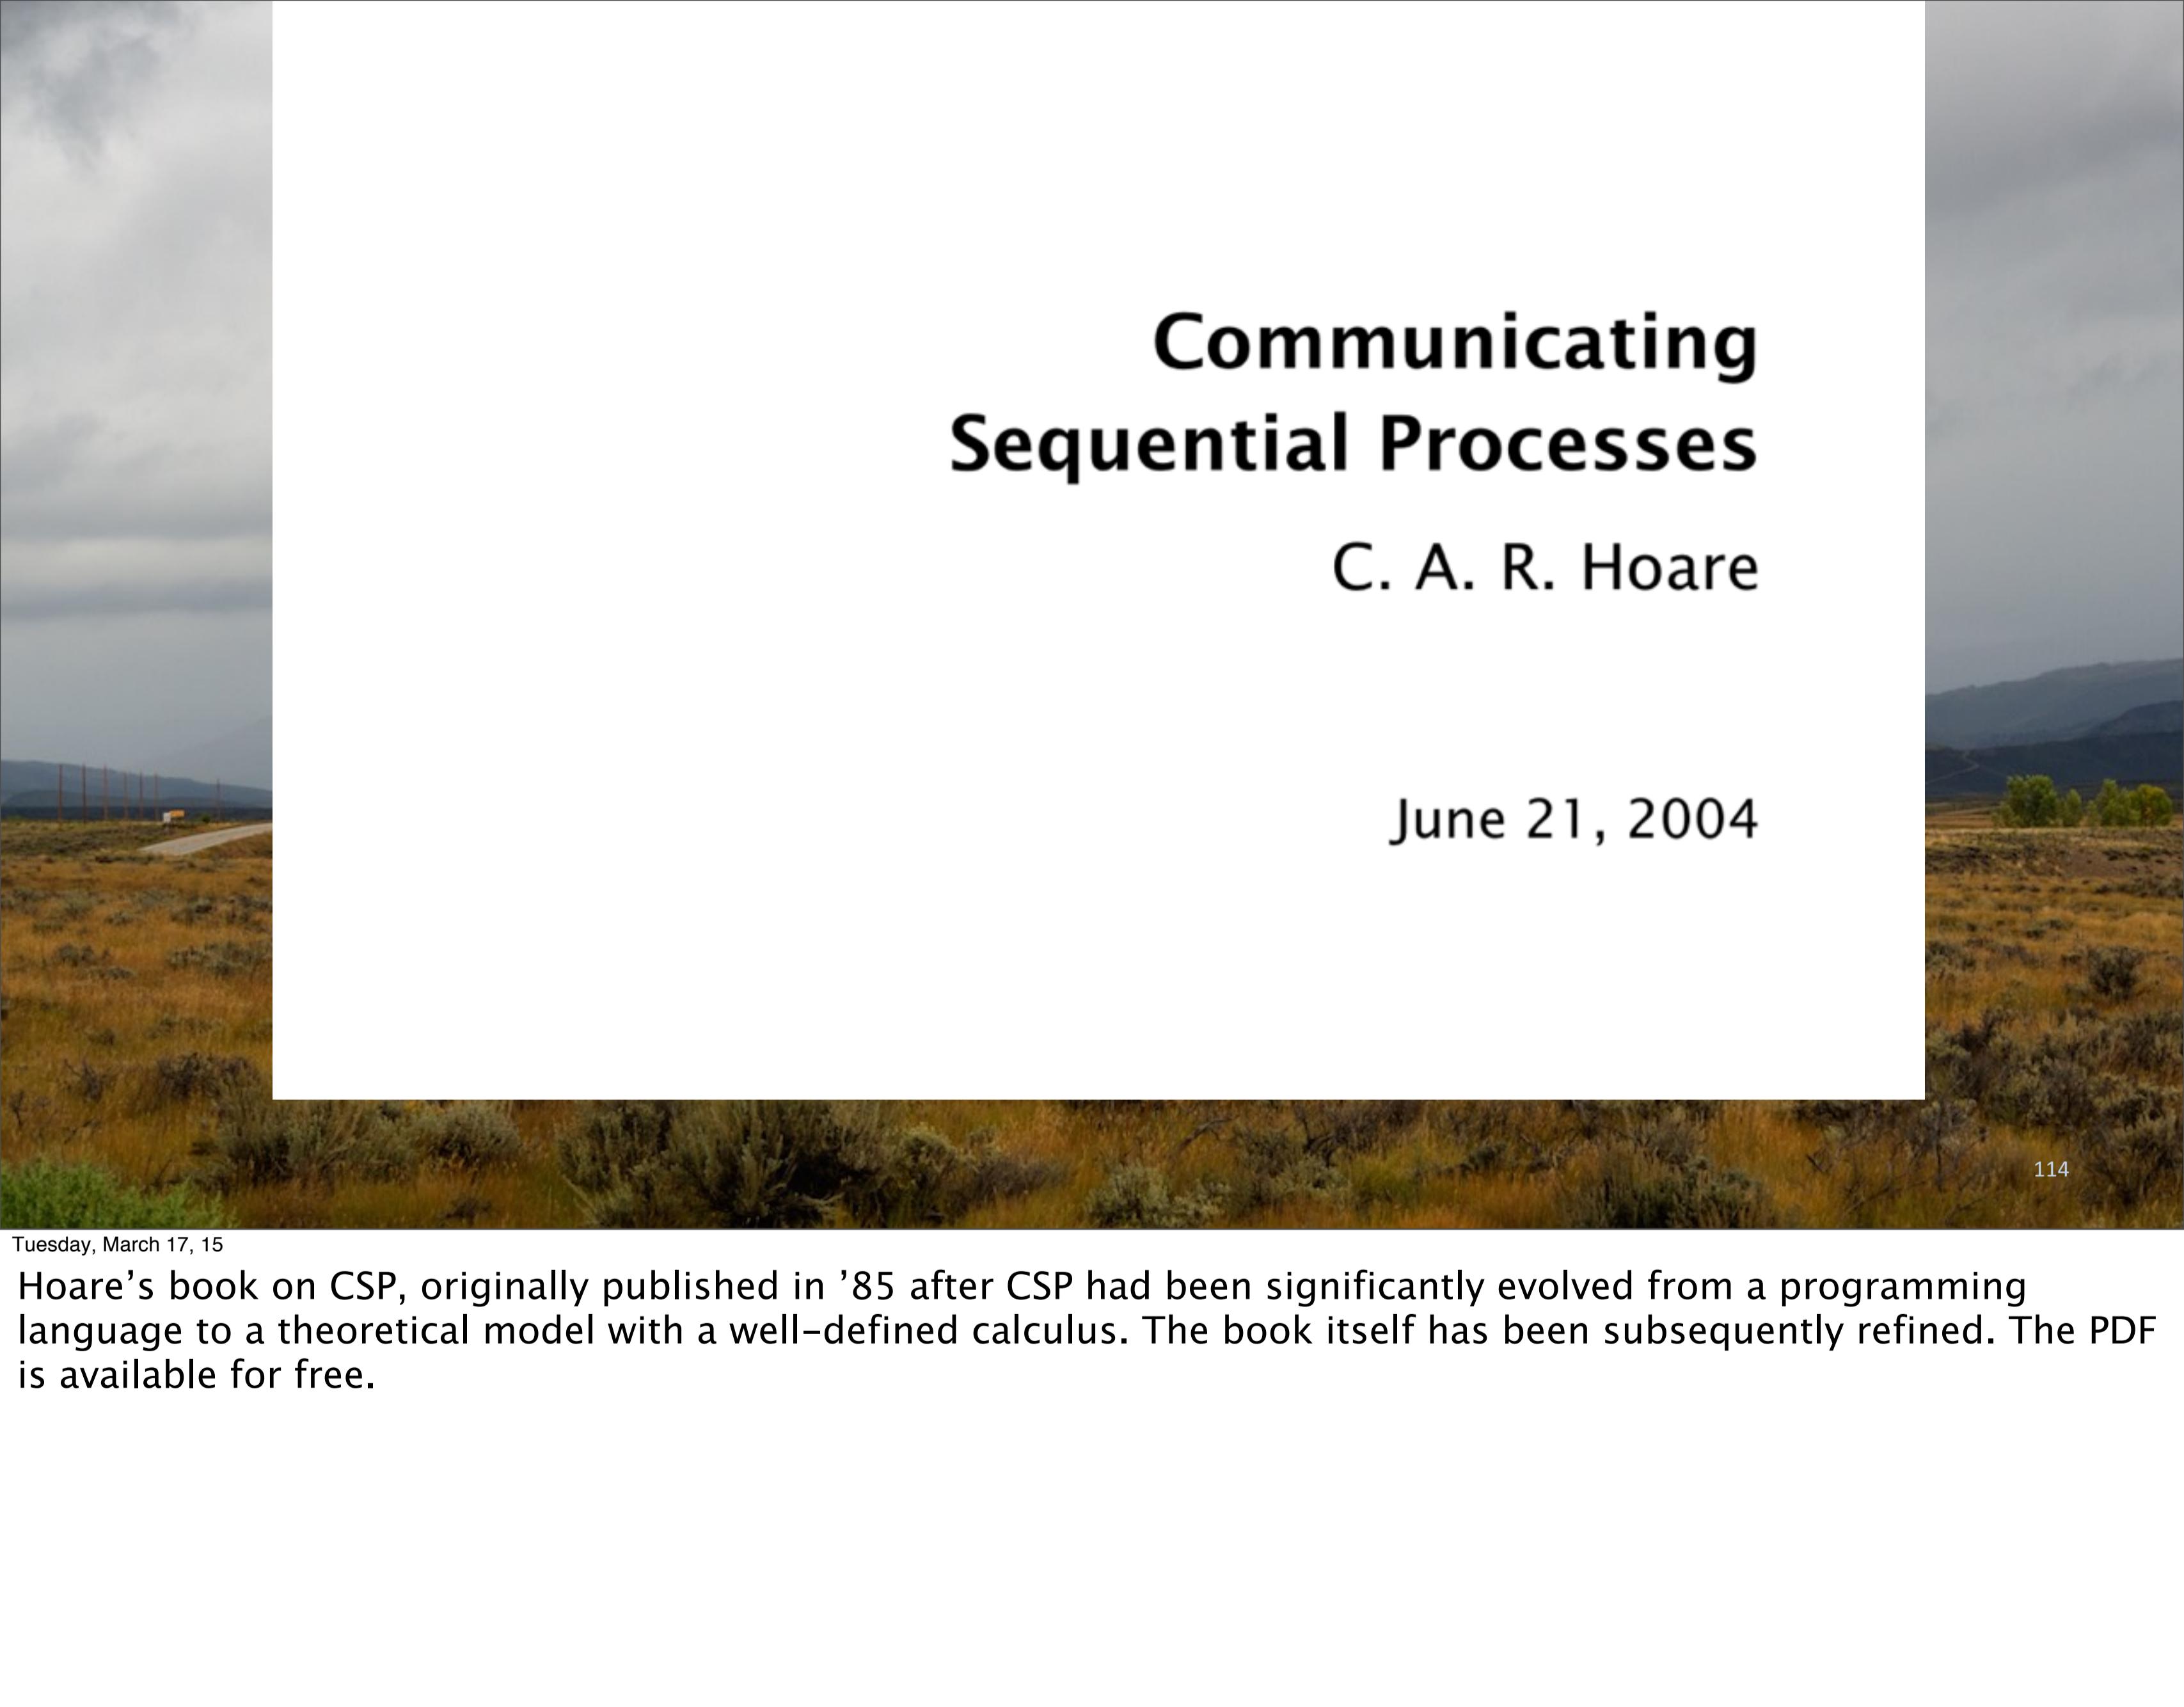
See also links earlier in the presentation.



113

Tuesday, March 17, 15

Lots of interesting practical ideas for combining functional programming and reactive approaches to class Domain-Driven Design by Eric Evans.



Communicating Sequential Processes

C. A. R. Hoare

June 21, 2004

114

Tuesday, March 17, 15

Hoare's book on CSP, originally published in '85 after CSP had been significantly evolved from a programming language to a theoretical model with a well-defined calculus. The book itself has been subsequently refined. The PDF is available for free.

The Theory and Practice of Concurrency

A.W. Roscoe

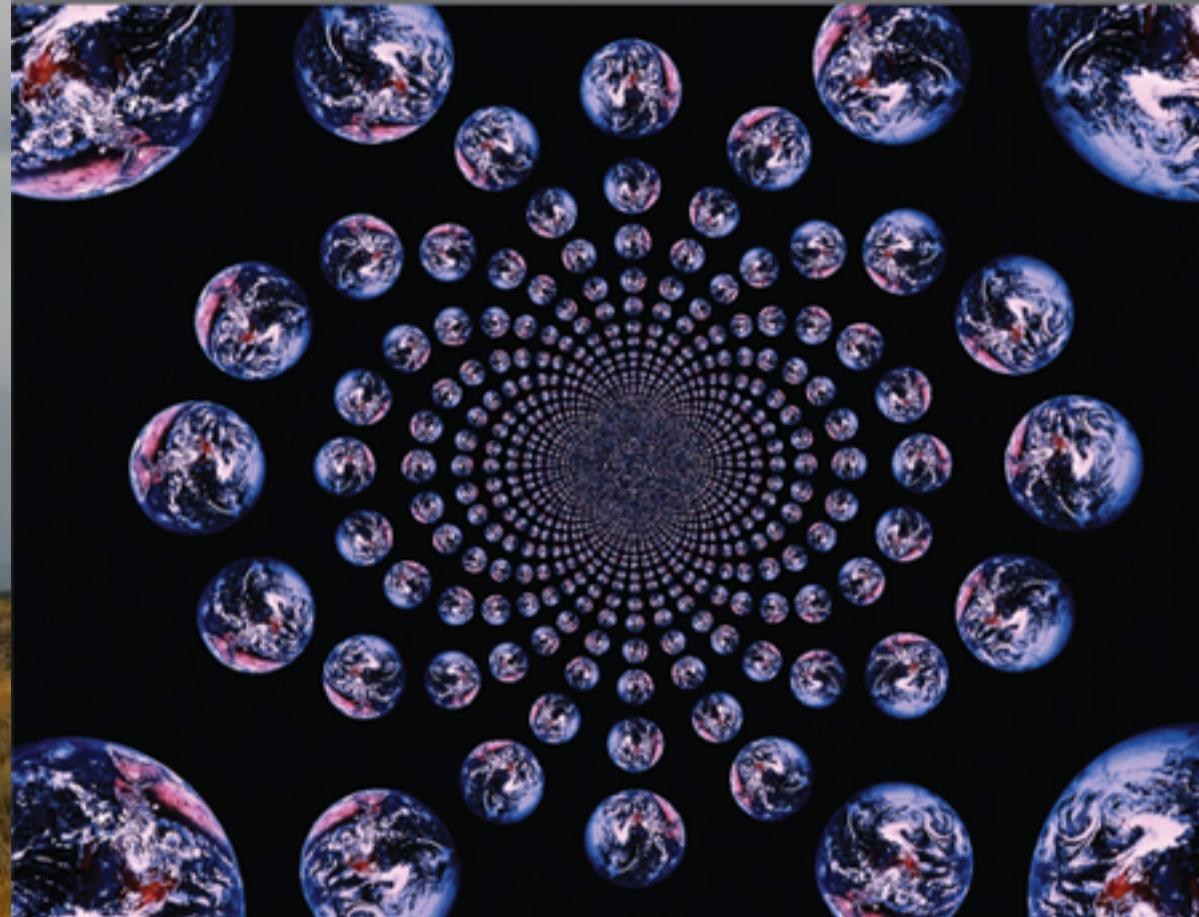
Published 1997, revised to 2000 and lightly revised to 2005.

The original version is in print in April 2005 with Prentice-Hall (Pearson).
This version is made available for personal reference only. This version is
copyright (©) Pearson and Bill Roscoe.

PROGRAMMING DISTRIBUTED COMPUTING SYSTEMS

A Foundational Approach

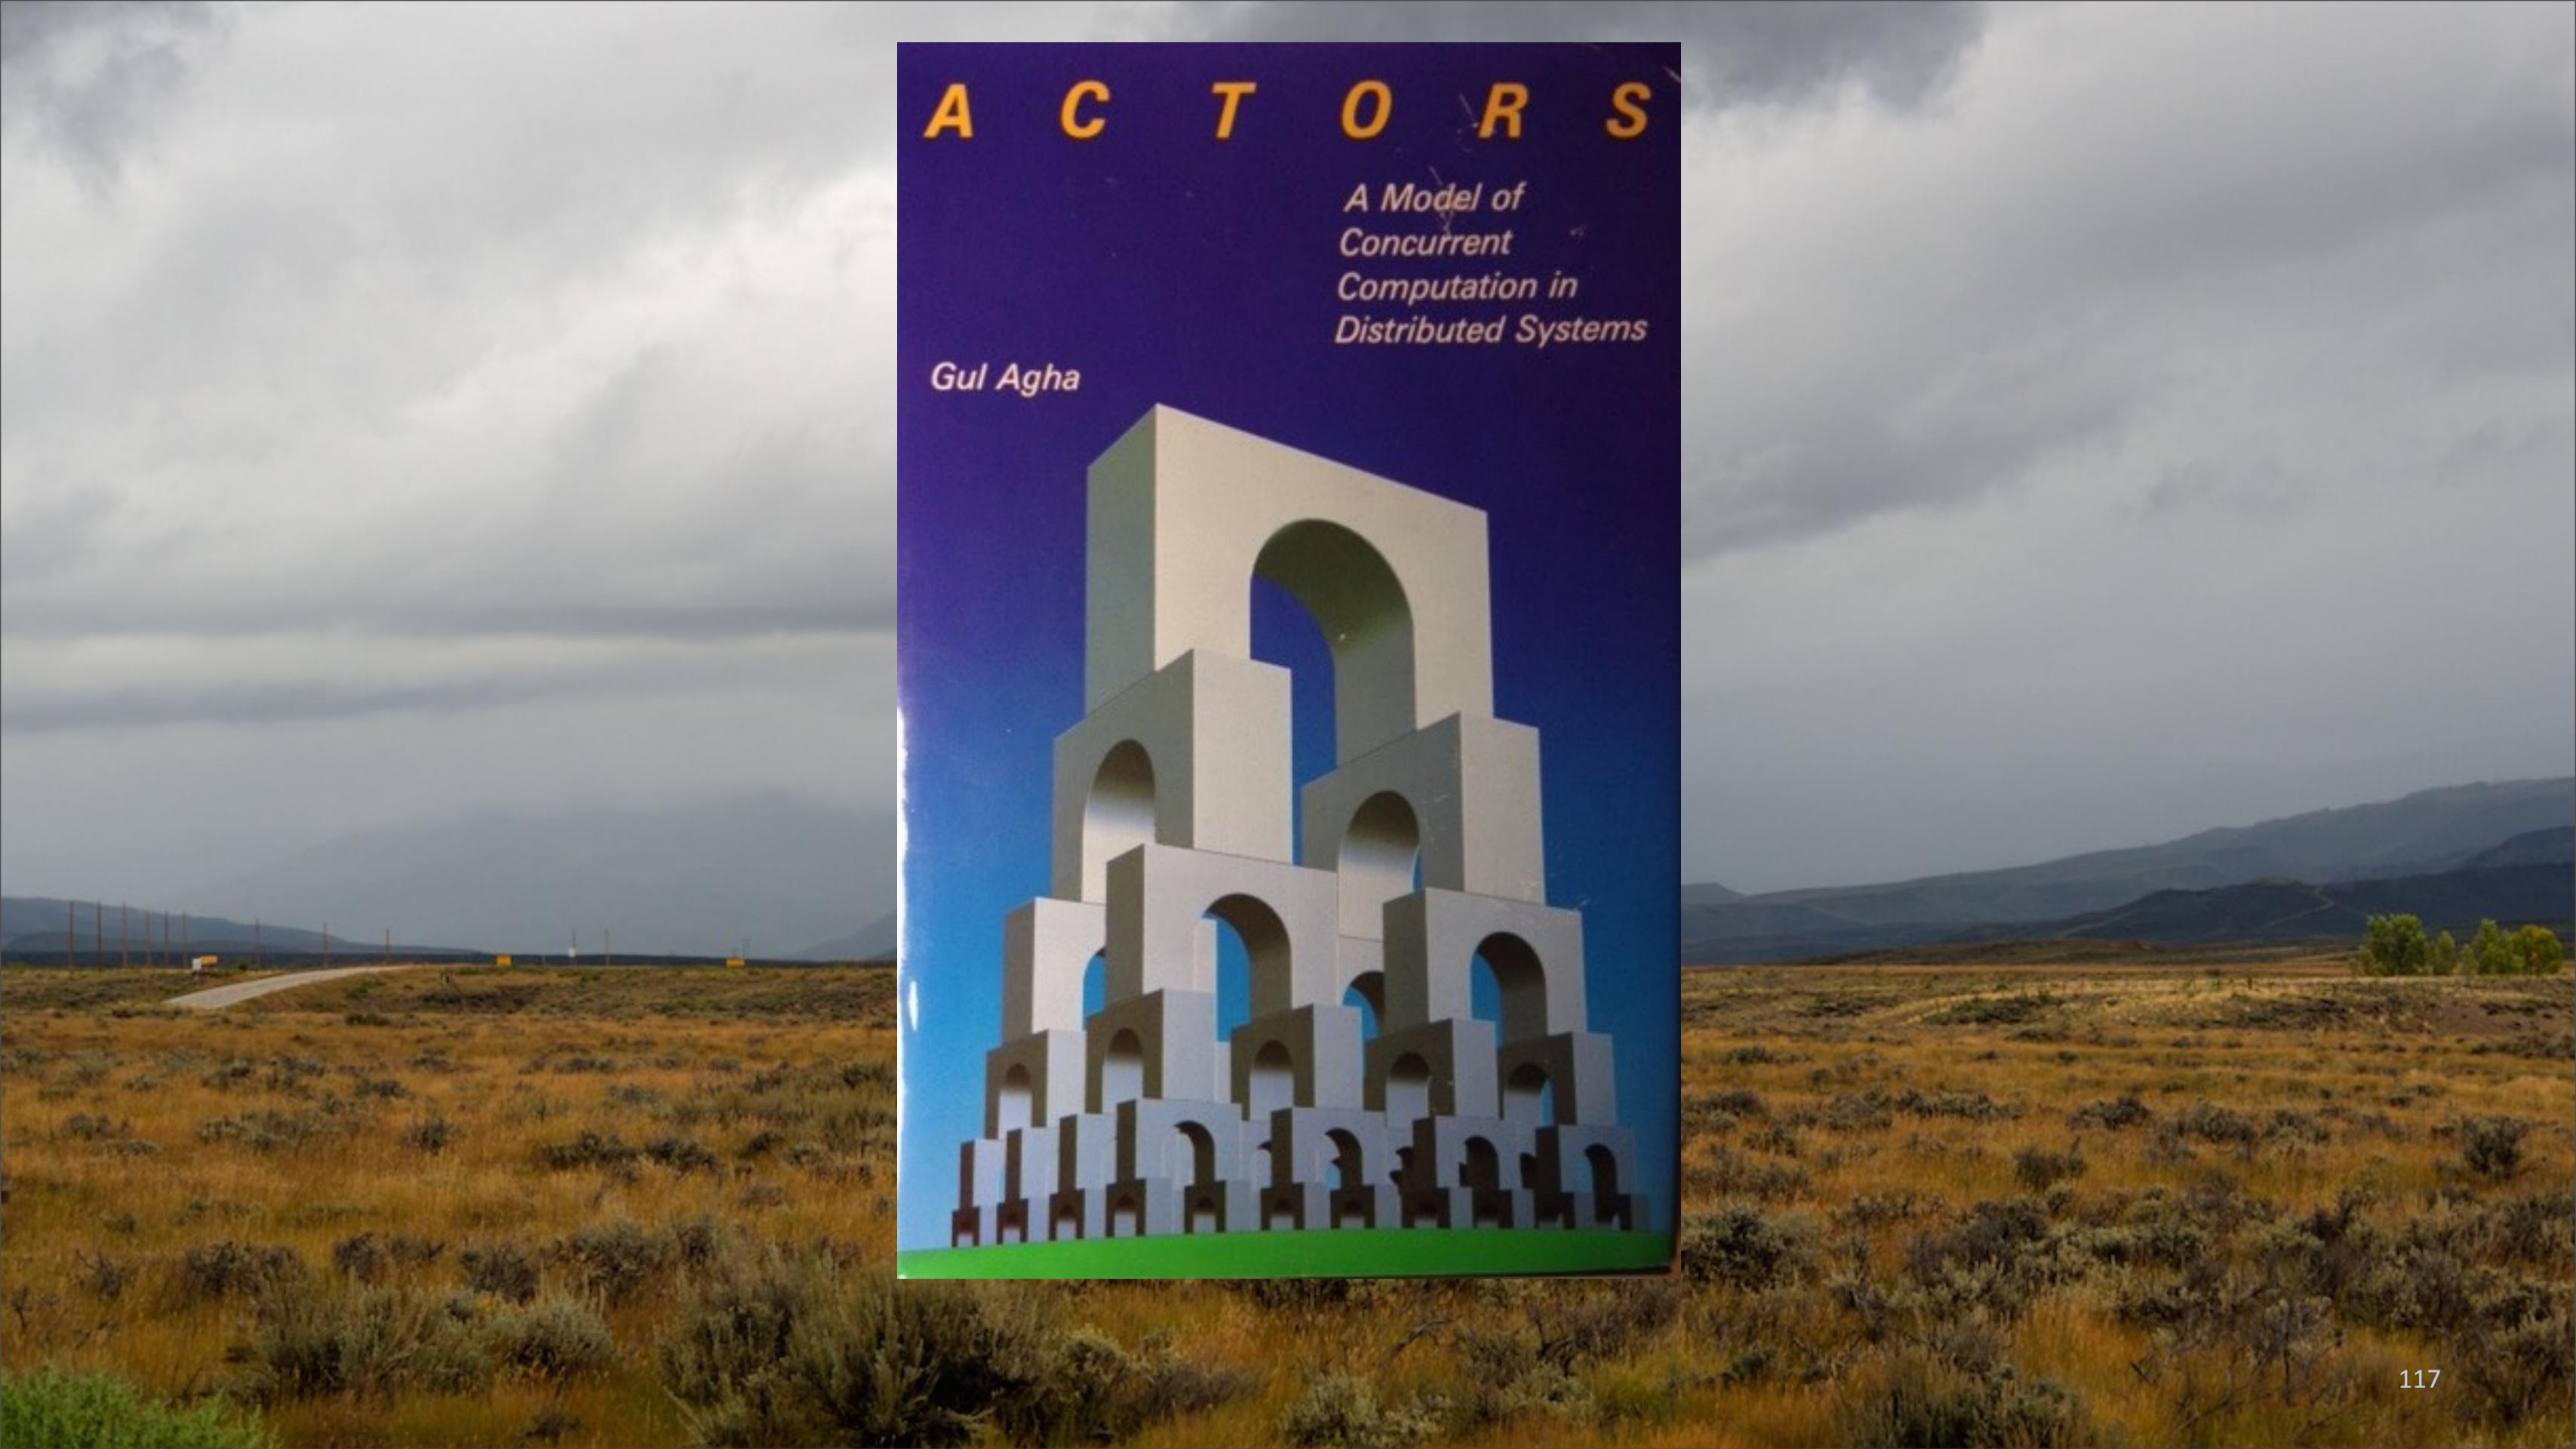
CARLOS A. VARELA



116

Tuesday, March 17, 15

A survey of theoretical models of distributed computing, starting with a summary of lambda calculus, then discussing the pi, join, and ambient calculi. Also discusses the actor model. The treatment is somewhat dry and could use more discussion of real-world implementations of these ideas, such as the Actor model in Erlang and Akka.



117

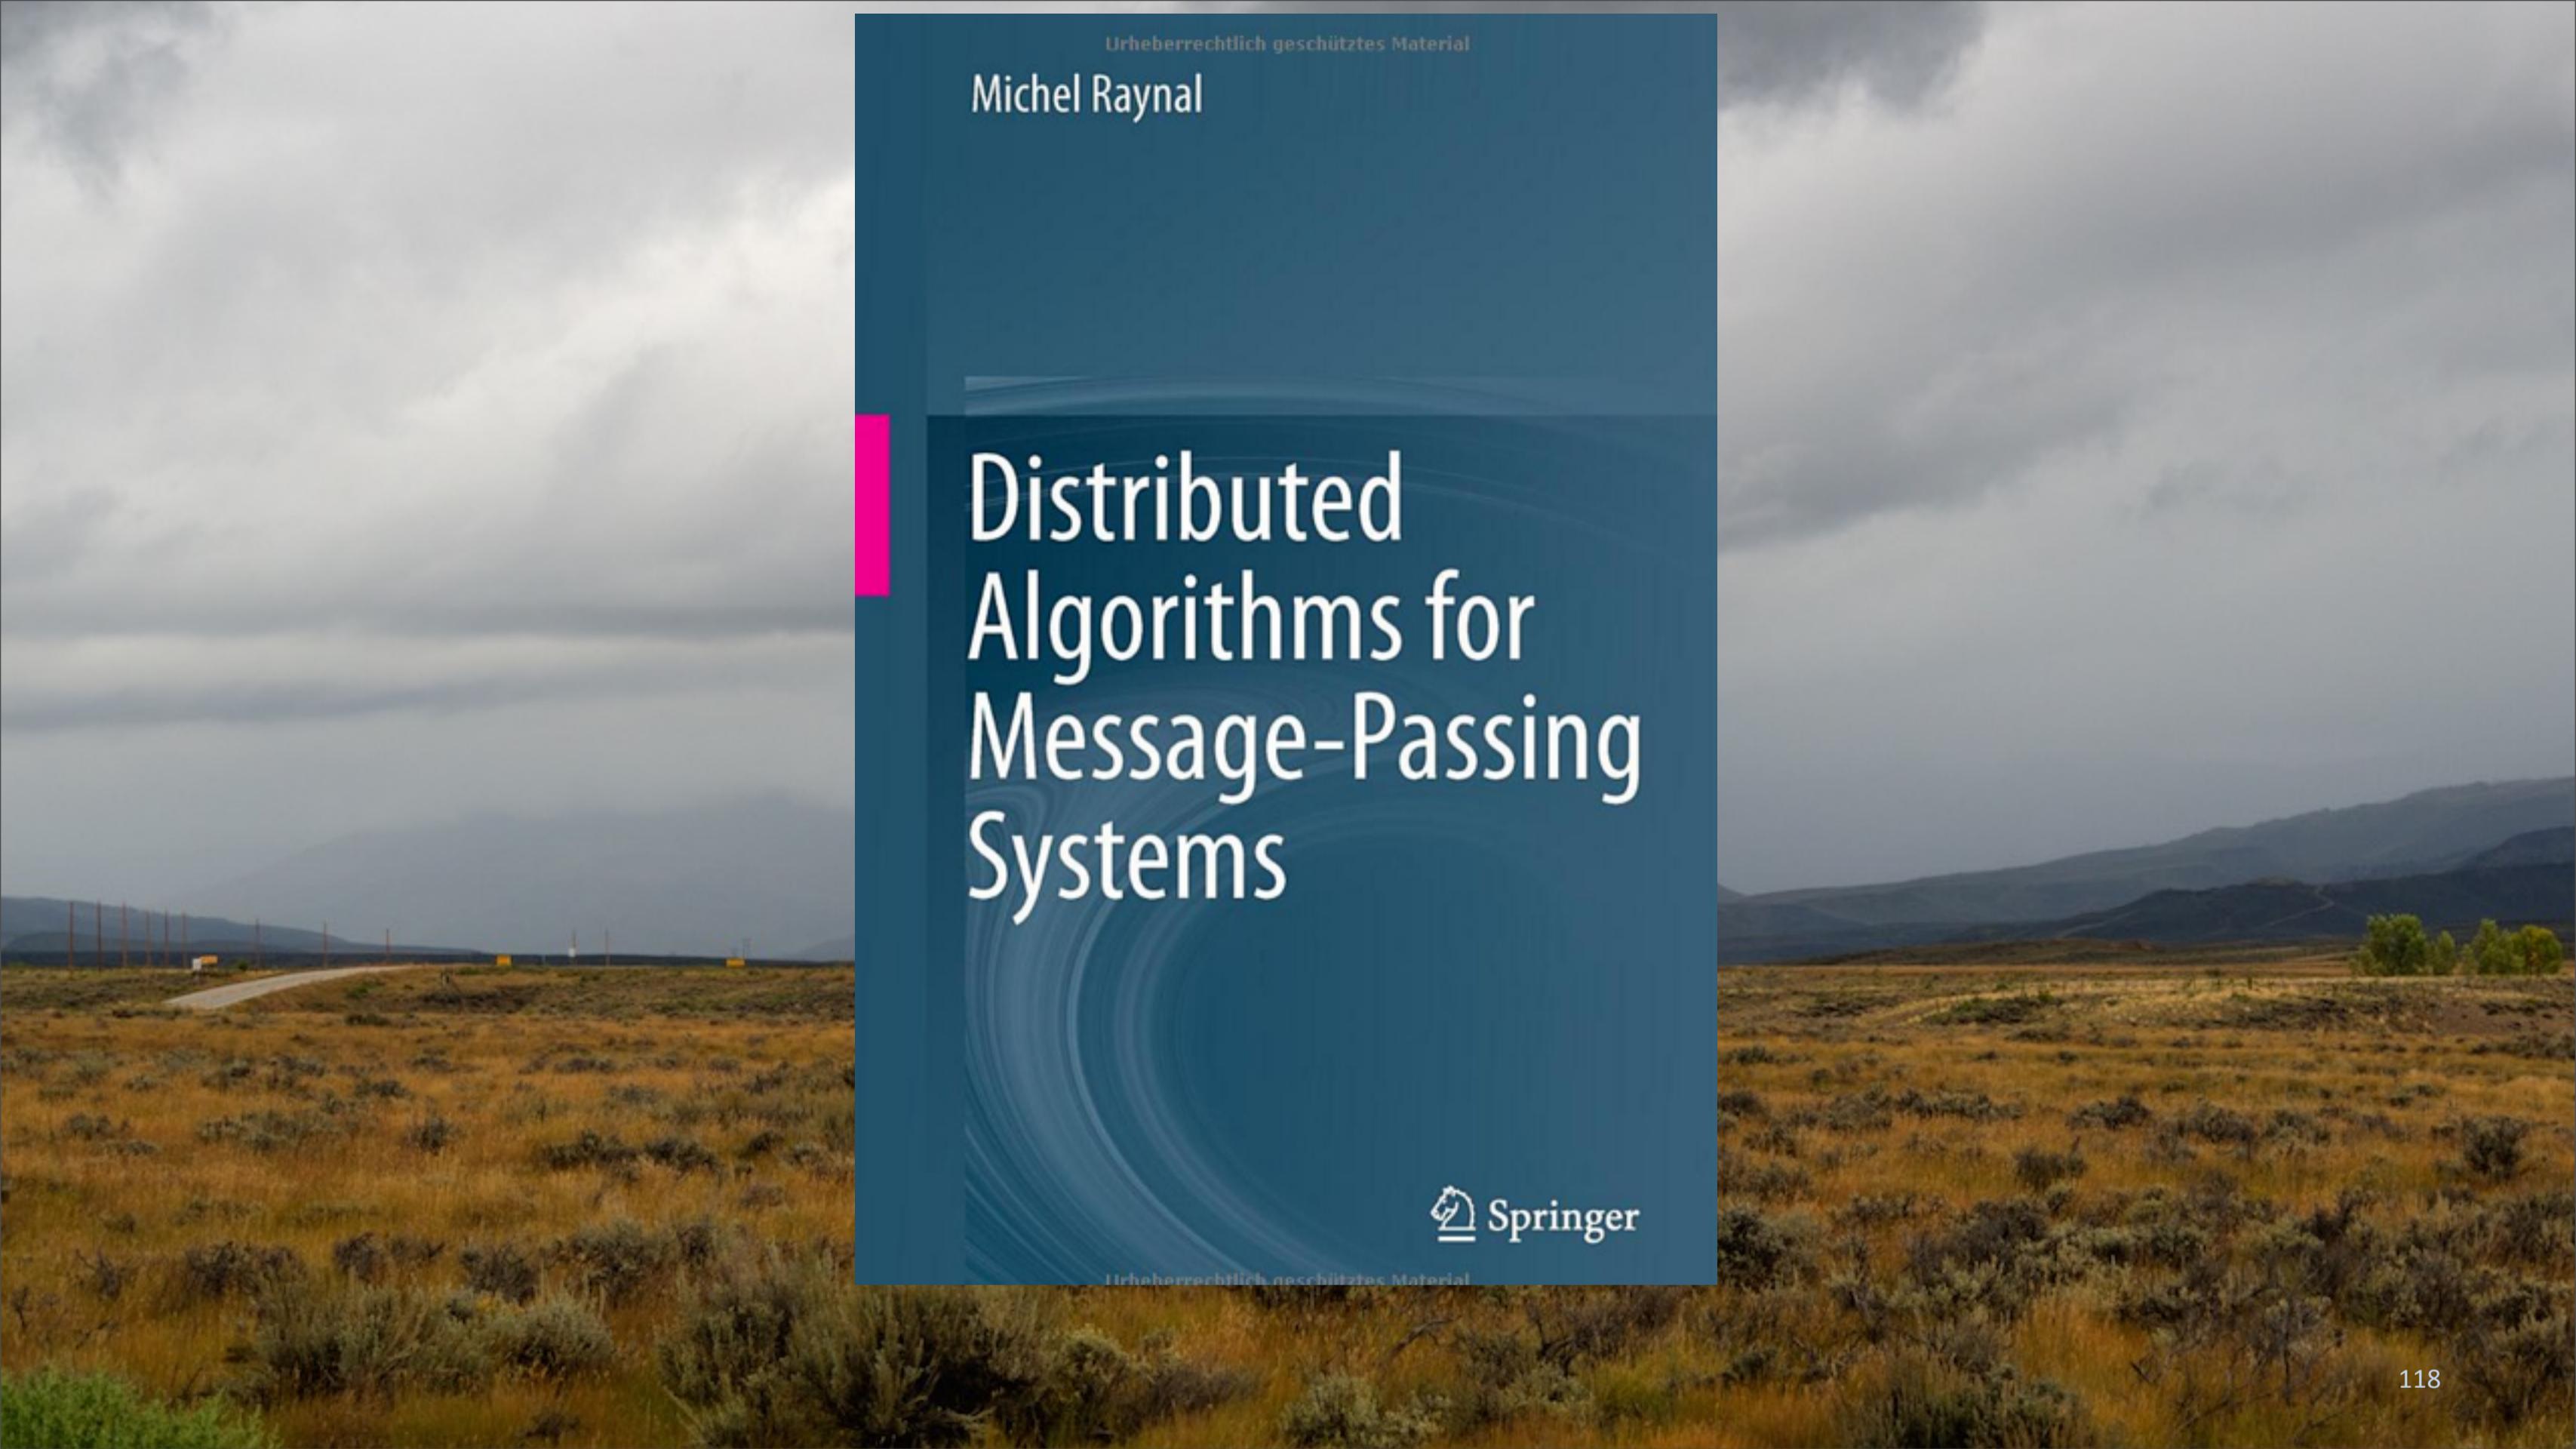
Tuesday, March 17, 15

Gul Agha was a grad student at MIT during the 80s and worked on the actor model with Hewitt and others. This book is based on his dissertation.

It doesn't discuss error handling, actor supervision, etc. as these concepts .

His thesis, <http://dspace.mit.edu/handle/1721.1/6952>, the basis for his book,<http://mitpress.mit.edu/books/actors>

See also Paper for a survey course with Rajesh Karmani, <http://www.cs.ucla.edu/~palsberg/course/cs239/papers/karmani-agha.pdf>



Tuesday, March 17, 15

118

Survey of the classic graph traversal algorithms, algorithms for detecting failures in a cluster, leader election, etc.

DISTRIBUTED ALGORITHMS

AN INTUITIVE APPROACH



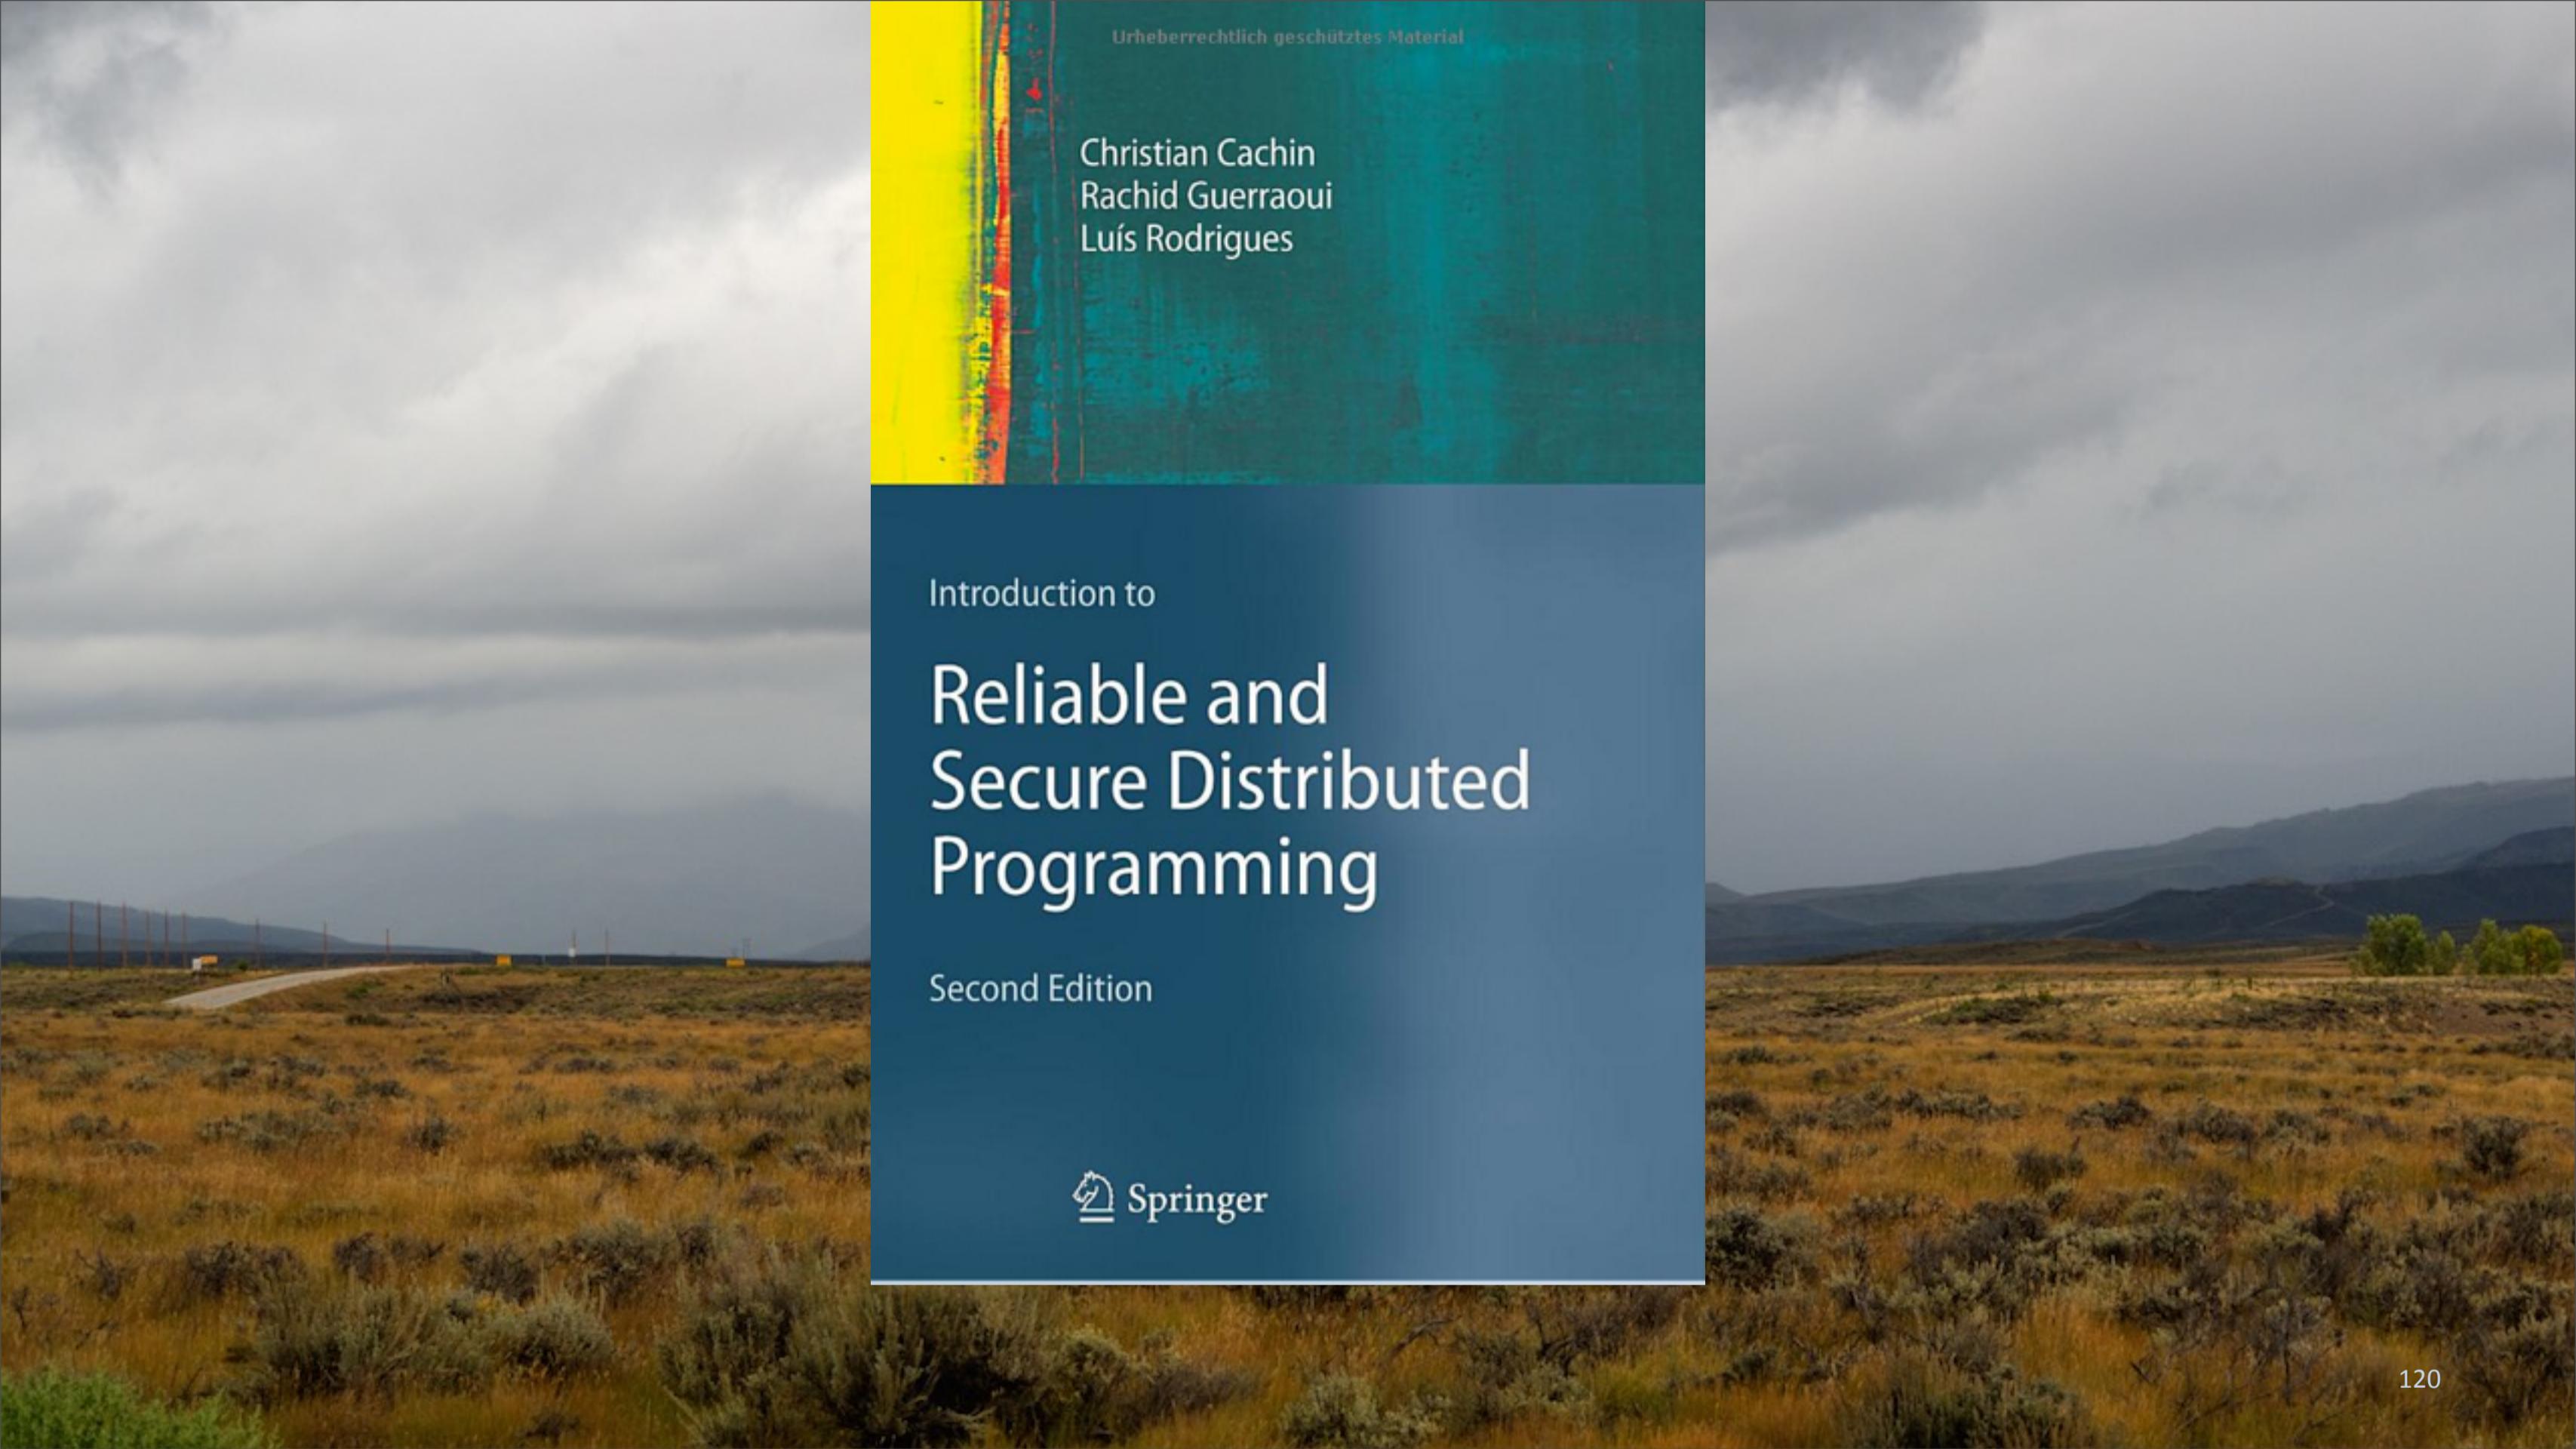
WAN FOKKINK

Copyrighted material

119

Tuesday, March 17, 15

A less comprehensive and formal, but more intuitive approach to fundamental algorithms.



120

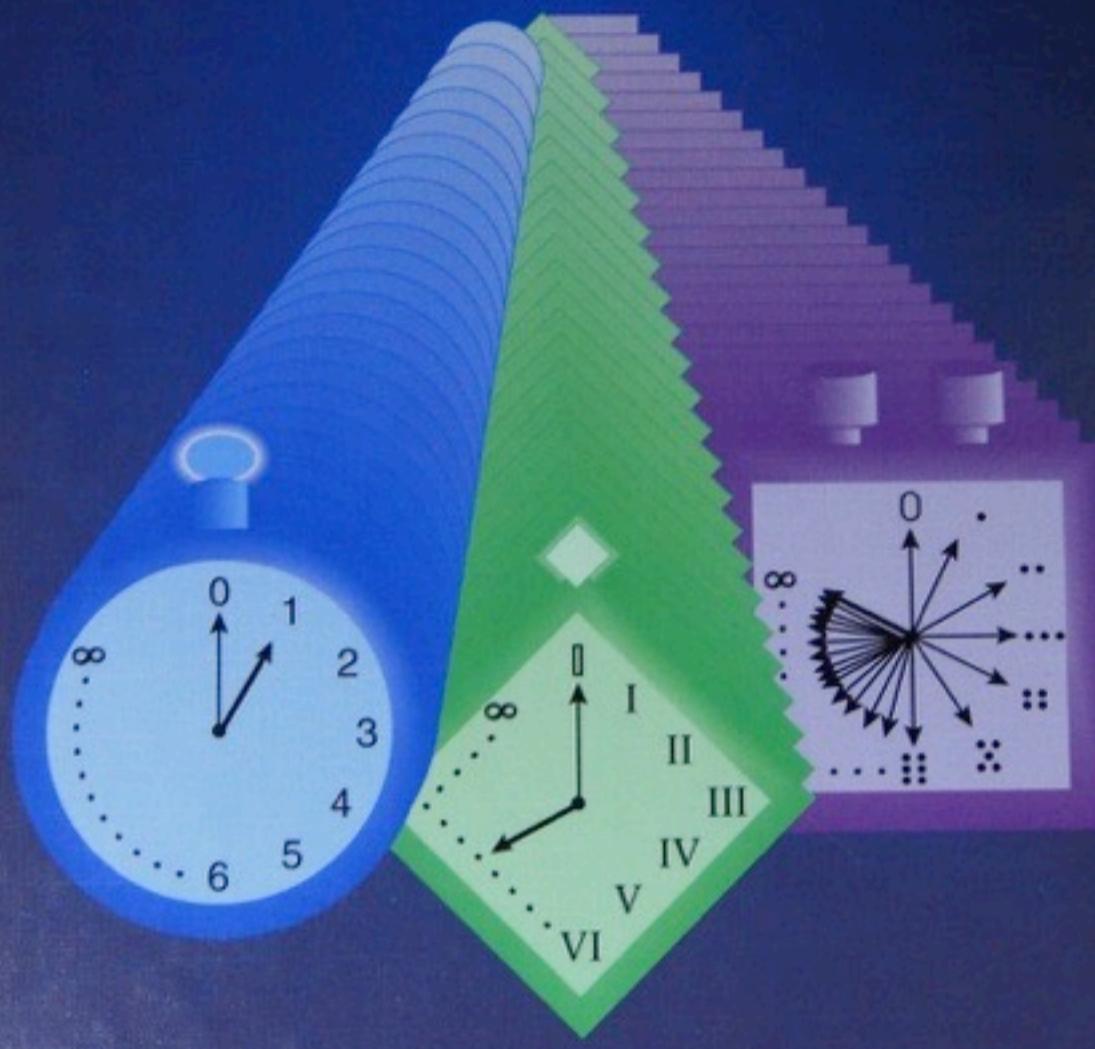
Tuesday, March 17, 15

Comprehensive and somewhat formal like Raynal's book, but more focused on modeling common failures in real systems.

Zohar Manna
Amir Pnueli

The Temporal Logic of Reactive and Concurrent Systems

•Specification•

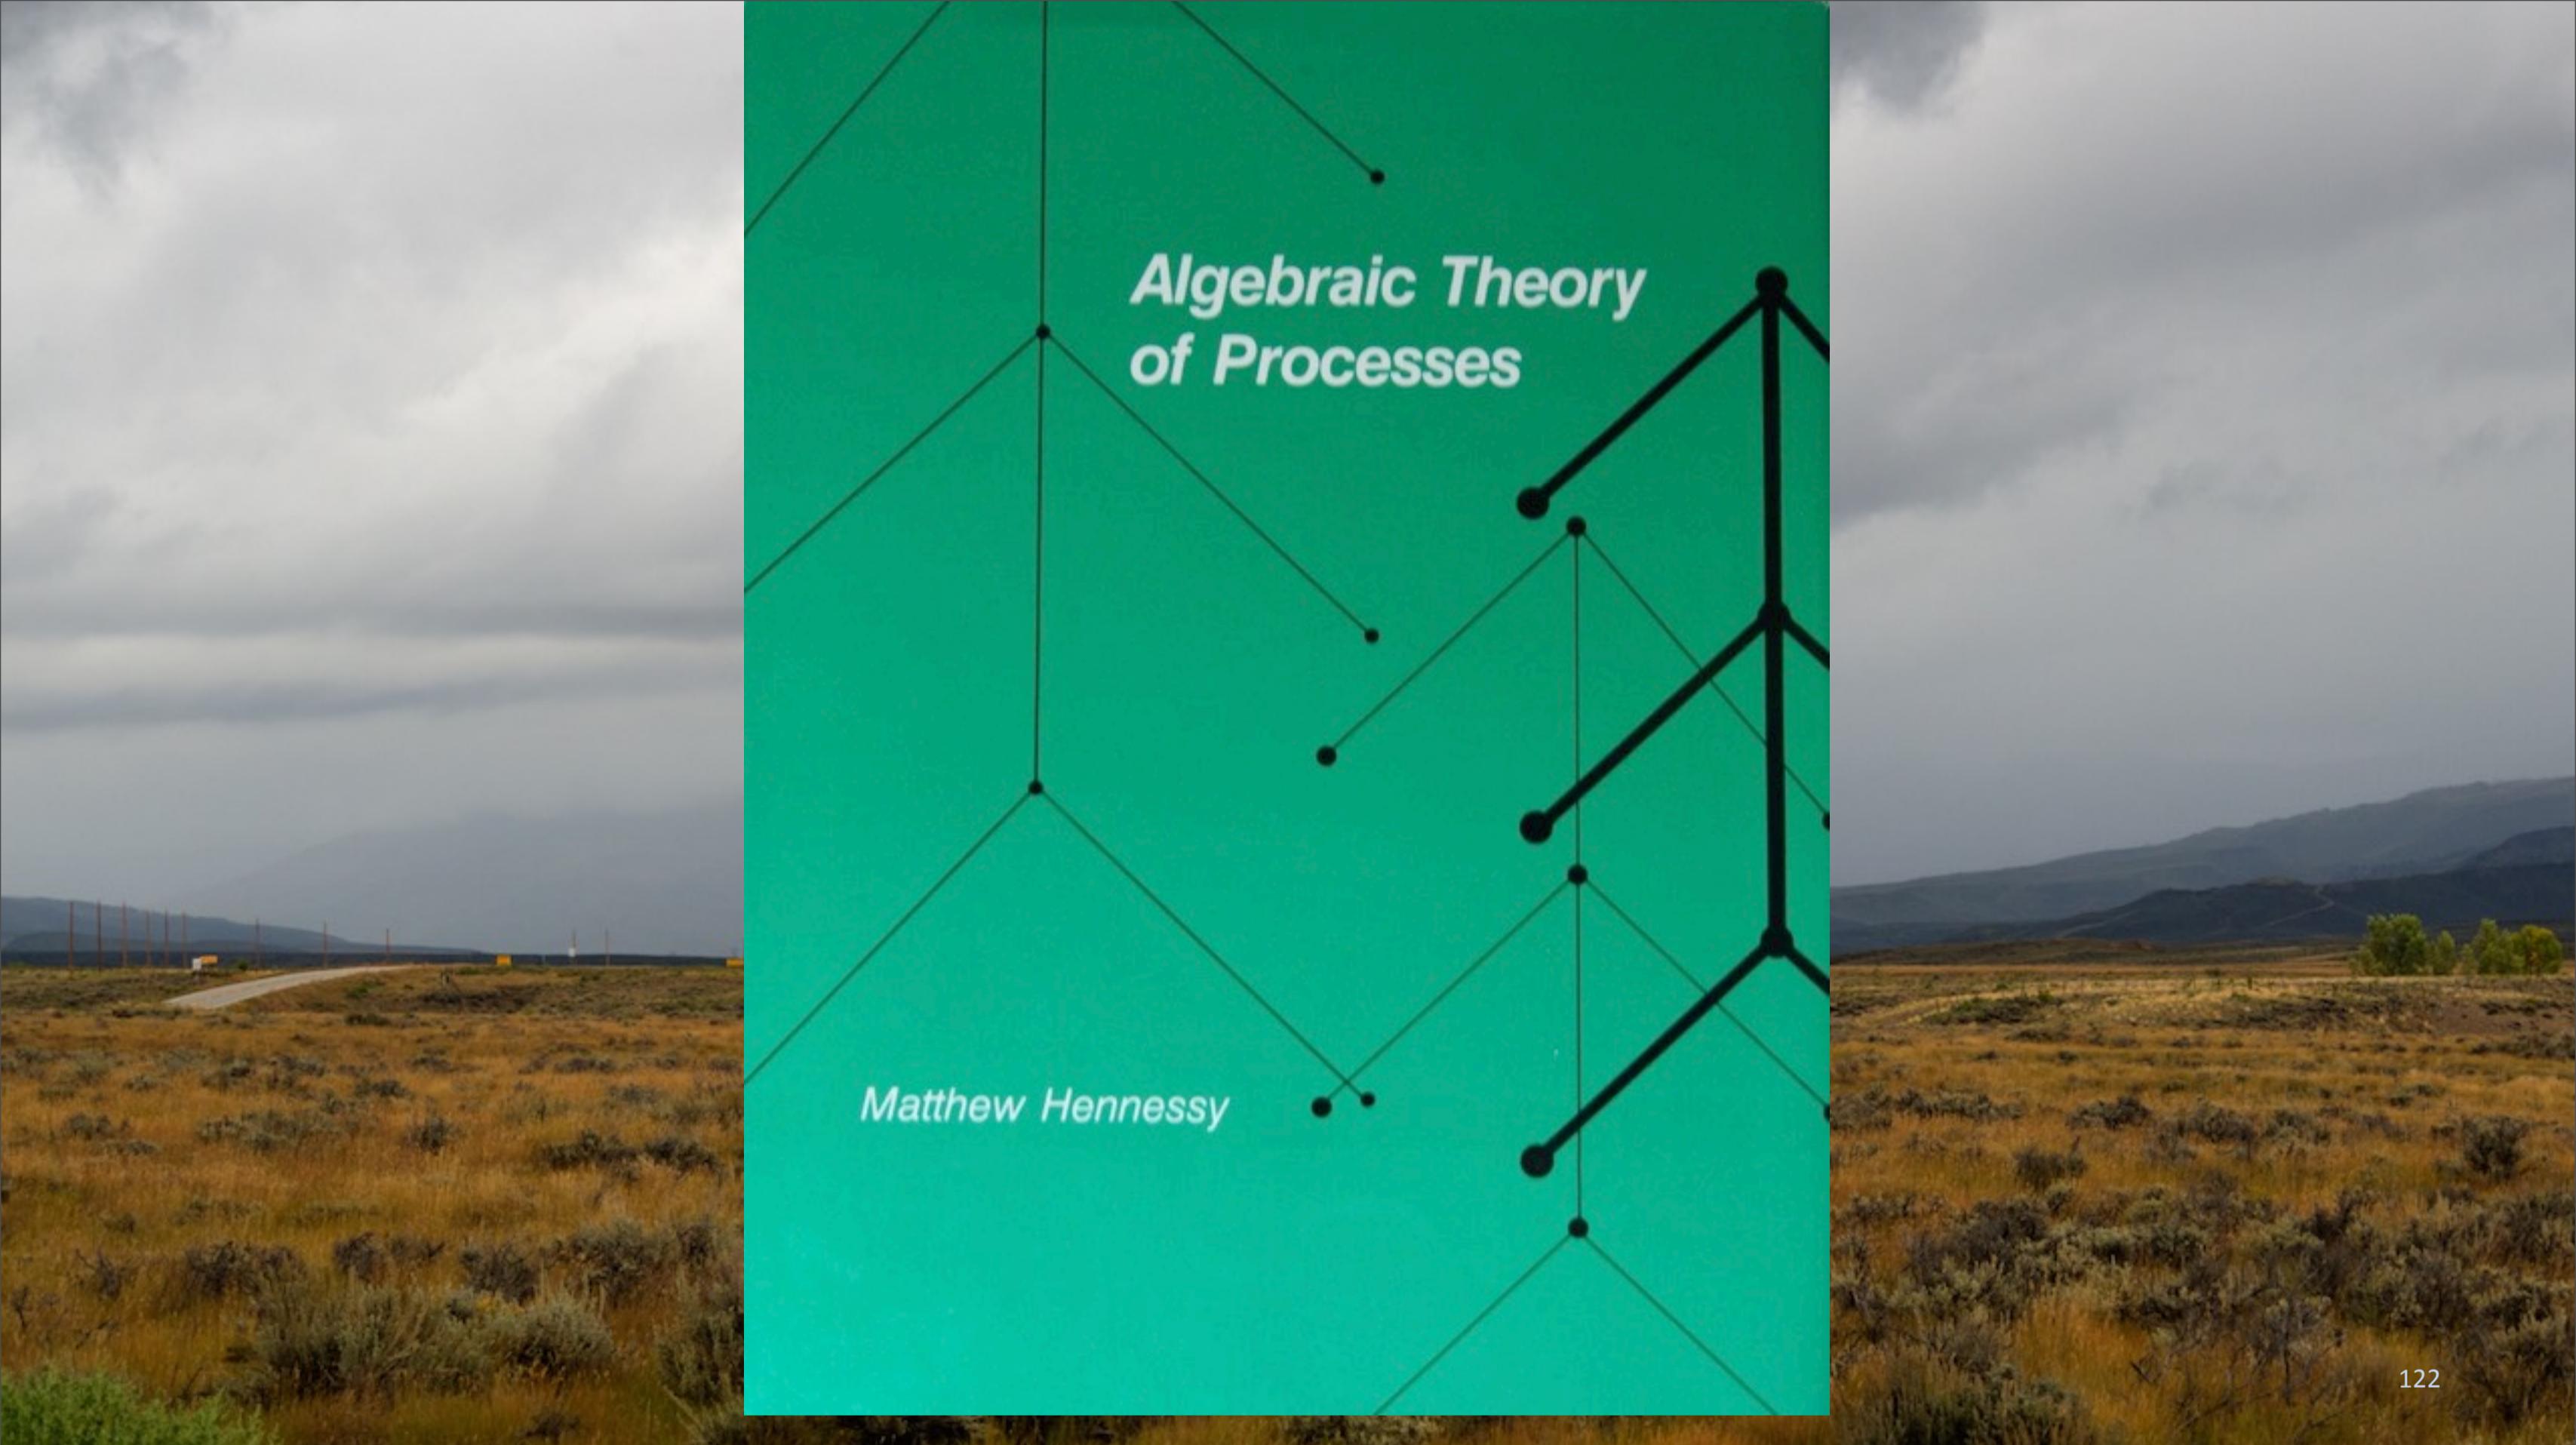


Springer-Verlag

121

Tuesday, March 17, 15

1992: Yes, “Reactive” isn’t new ;) This book is lays out a theoretical model for specifying and proving “reactive” concurrent systems based on temporal logic. While its goal is to prevent logic errors, It doesn’t discuss handling failures from environmental or other external causes in great depth.



122

Tuesday, March 17, 15

1988: Another treatment of concurrency using algebra. It's not based on CSP, but it has similar constructs.

DISTRIBUTED COMPUTING *through* COMBINATORIAL TOPOLOGY



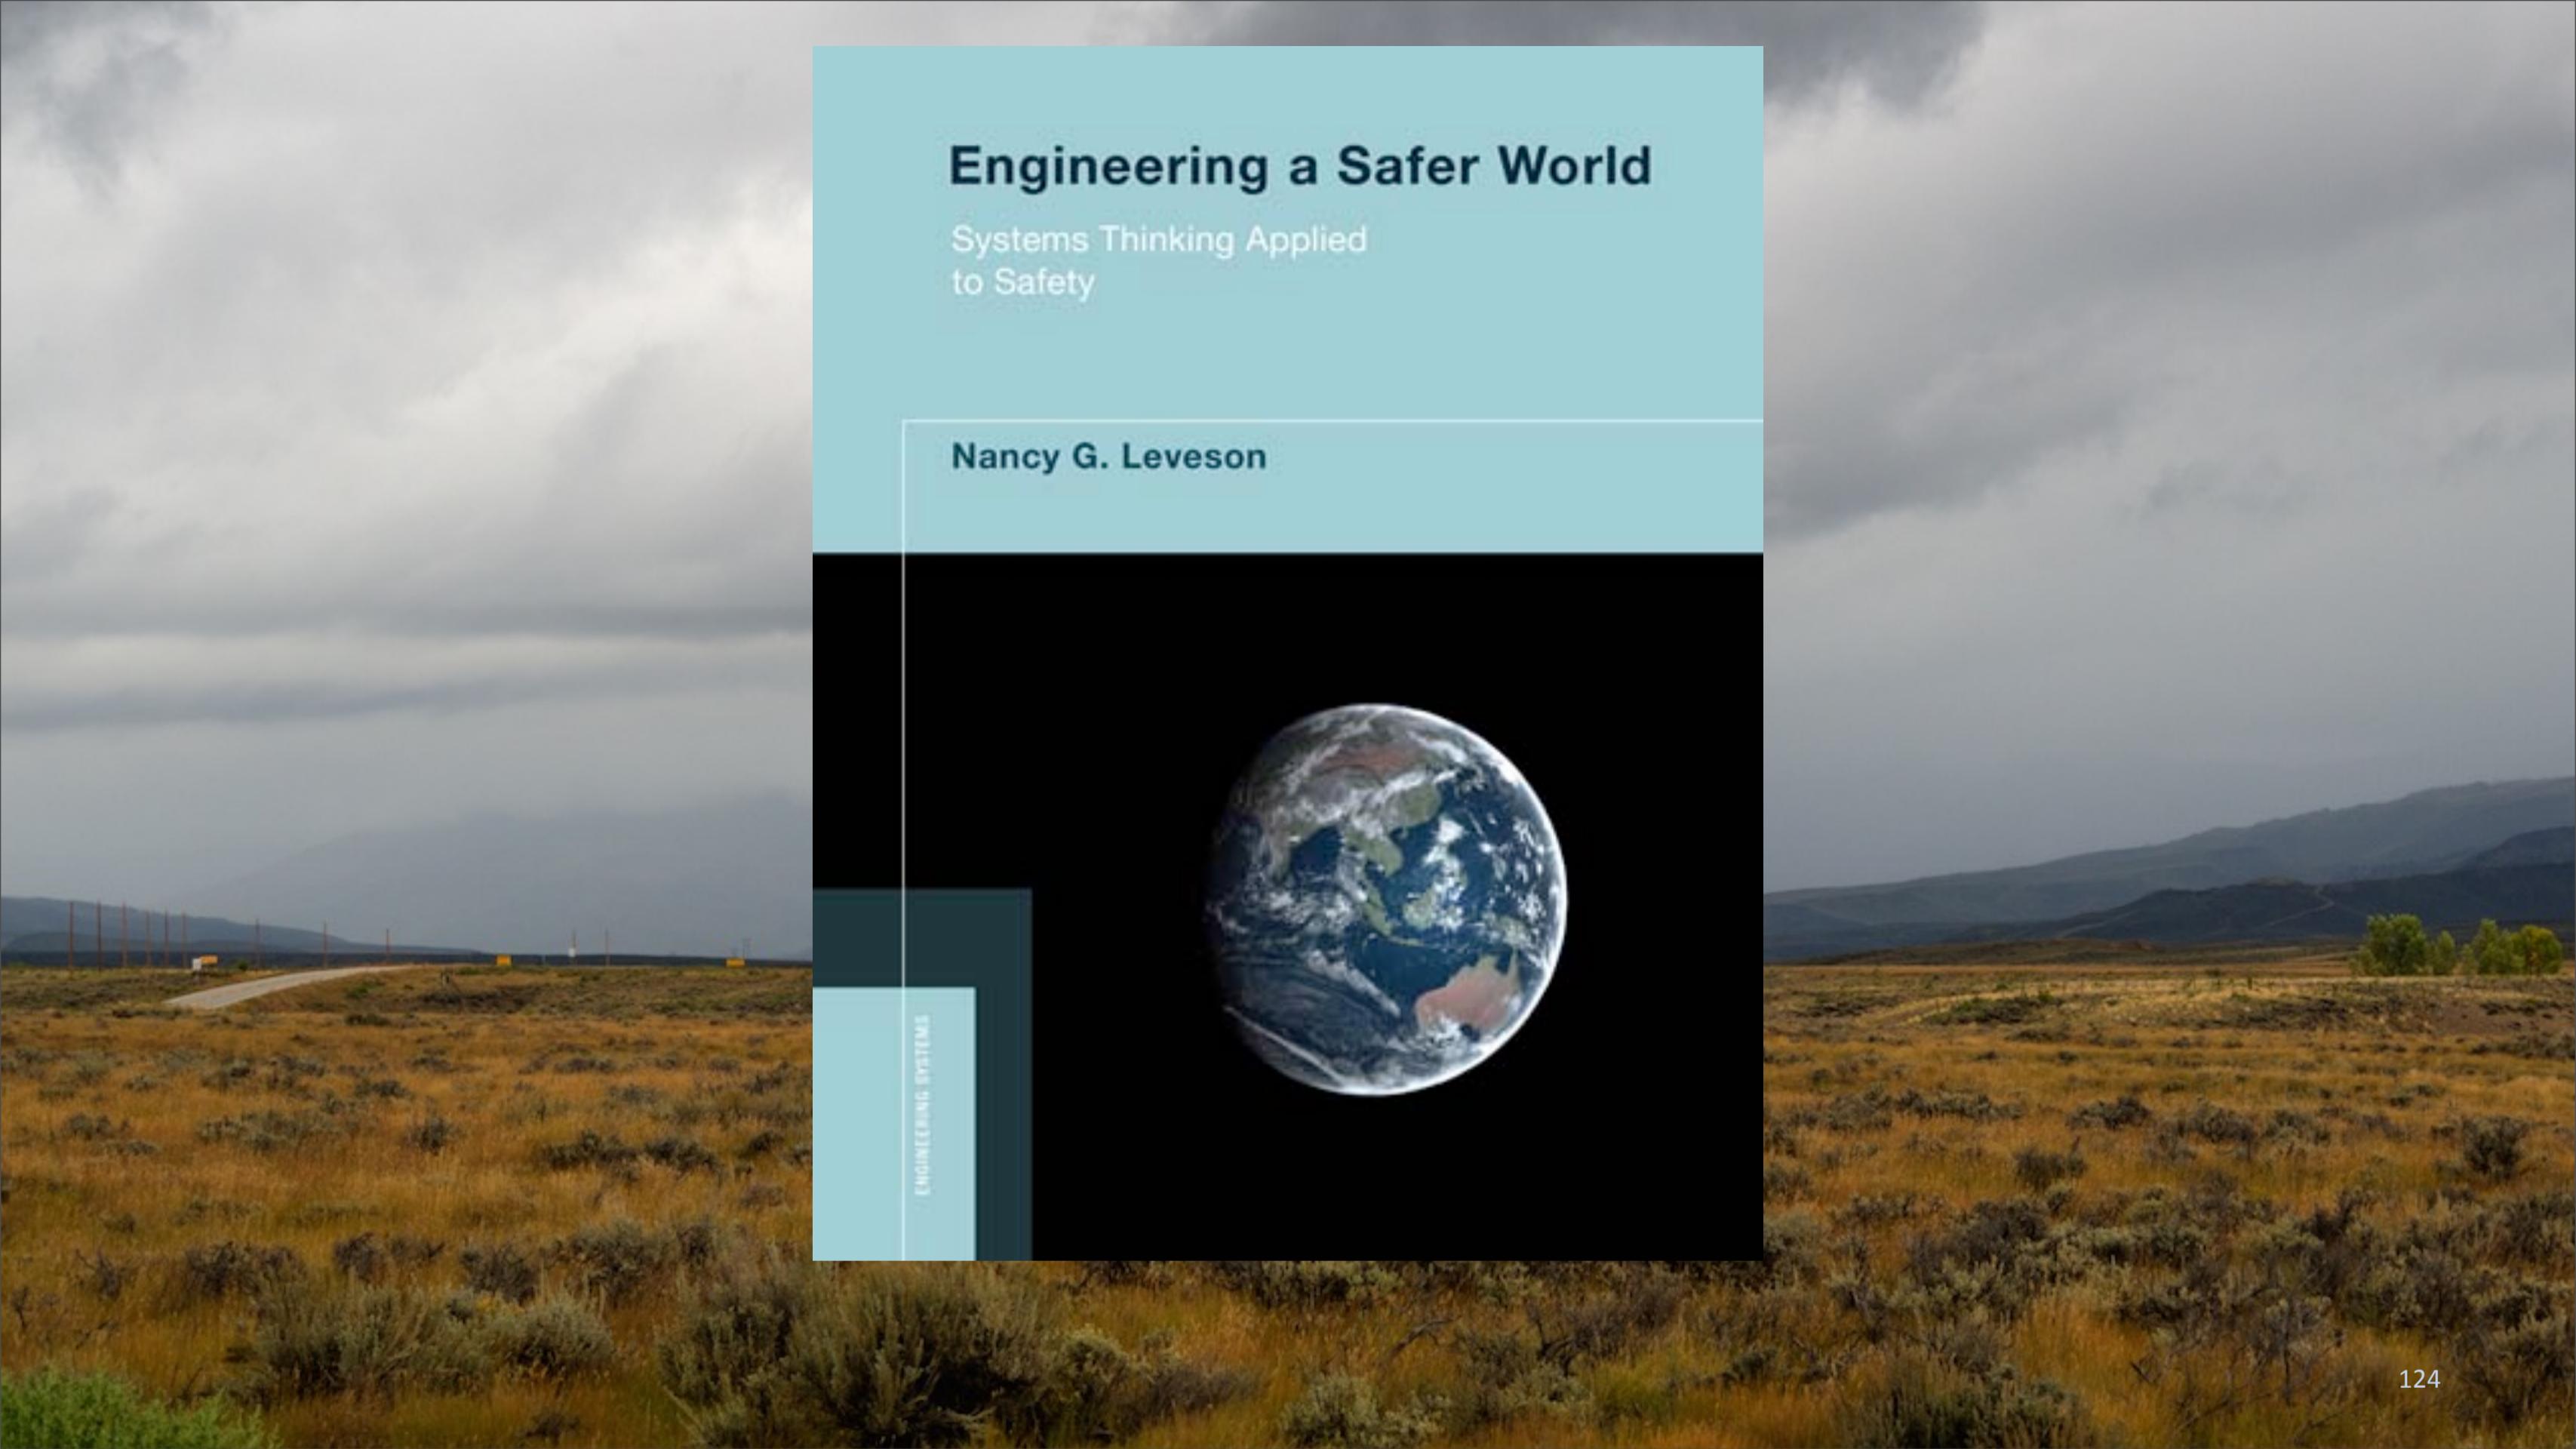
Copyrighted Material

Maurice Herlihy
Dmitry Kozlov
Sergio Rajsbaum

123

Tuesday, March 17, 15

A recent text that applies combinatorics (counting things) and topology (properties of geometric shapes) to the analysis of distributed systems. Aims to be pragmatic for real-world scenarios, like networks and other physical systems where failures are practical concerns.



Engineering a Safer World

Systems Thinking Applied
to Safety

Nancy G. Leveson

ENGINEERING SYSTEMS



Tuesday, March 17, 15

<http://mitpress.mit.edu/books/engineering-safer-world>

Farther afield, this book discusses safety concerns from a systems engineering perspective.