

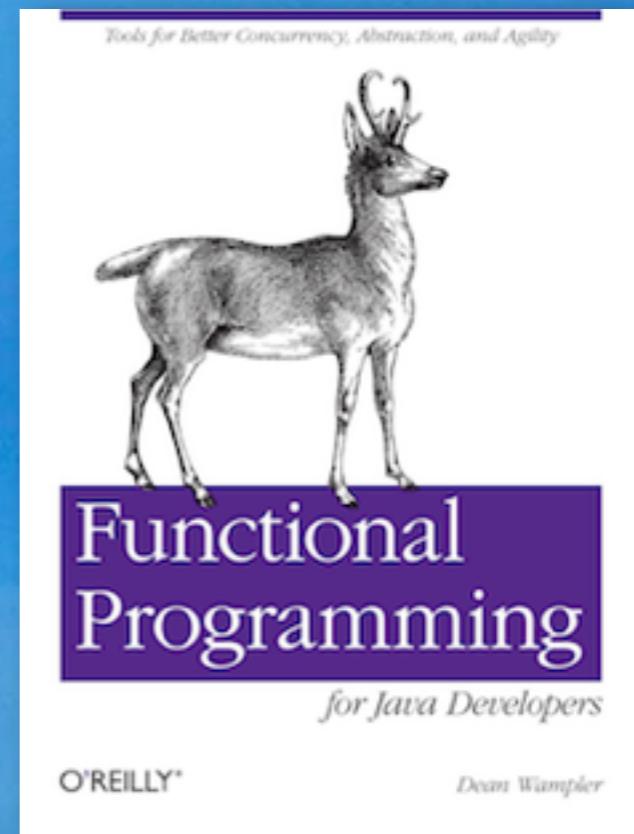
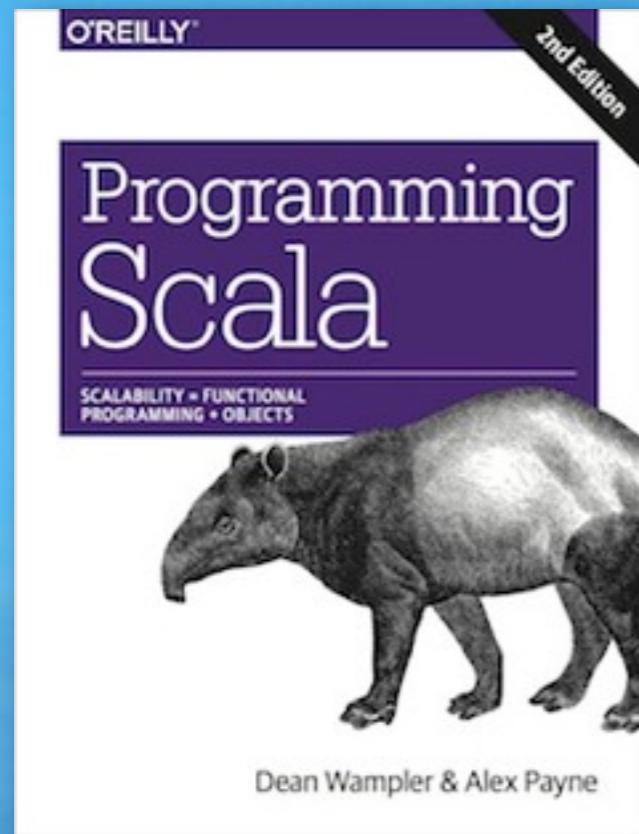


 **Typesafe**

Philly ETE 2014
April 22-23, 2014
@deanwampler
polyglotprogramming.com/talks

Why Spark Is
the Next Top
(Compute)
Model

Dean Wampler



dean.wampler@typesafe.com
polyglotprogramming.com/talks
@deanwampler

Wednesday, April 23, 14

About me. You can find this presentation and others on Big Data and Scala at polyglotprogramming.com.
Programming Scala, 2nd Edition is forthcoming.
photo: Dusk at 30,000 ft above the Central Plains of the U.S. on a Winter's Day.

Or
this?

THE
Compleat Troller,
OR,
THE ART
OF
TROLLING.
WITH
A Description of all the Utensils,
Instruments, Tackling, and Mate-
rials requisite thereto : With Rules
and Directions how to use them

Hadoop circa 2013

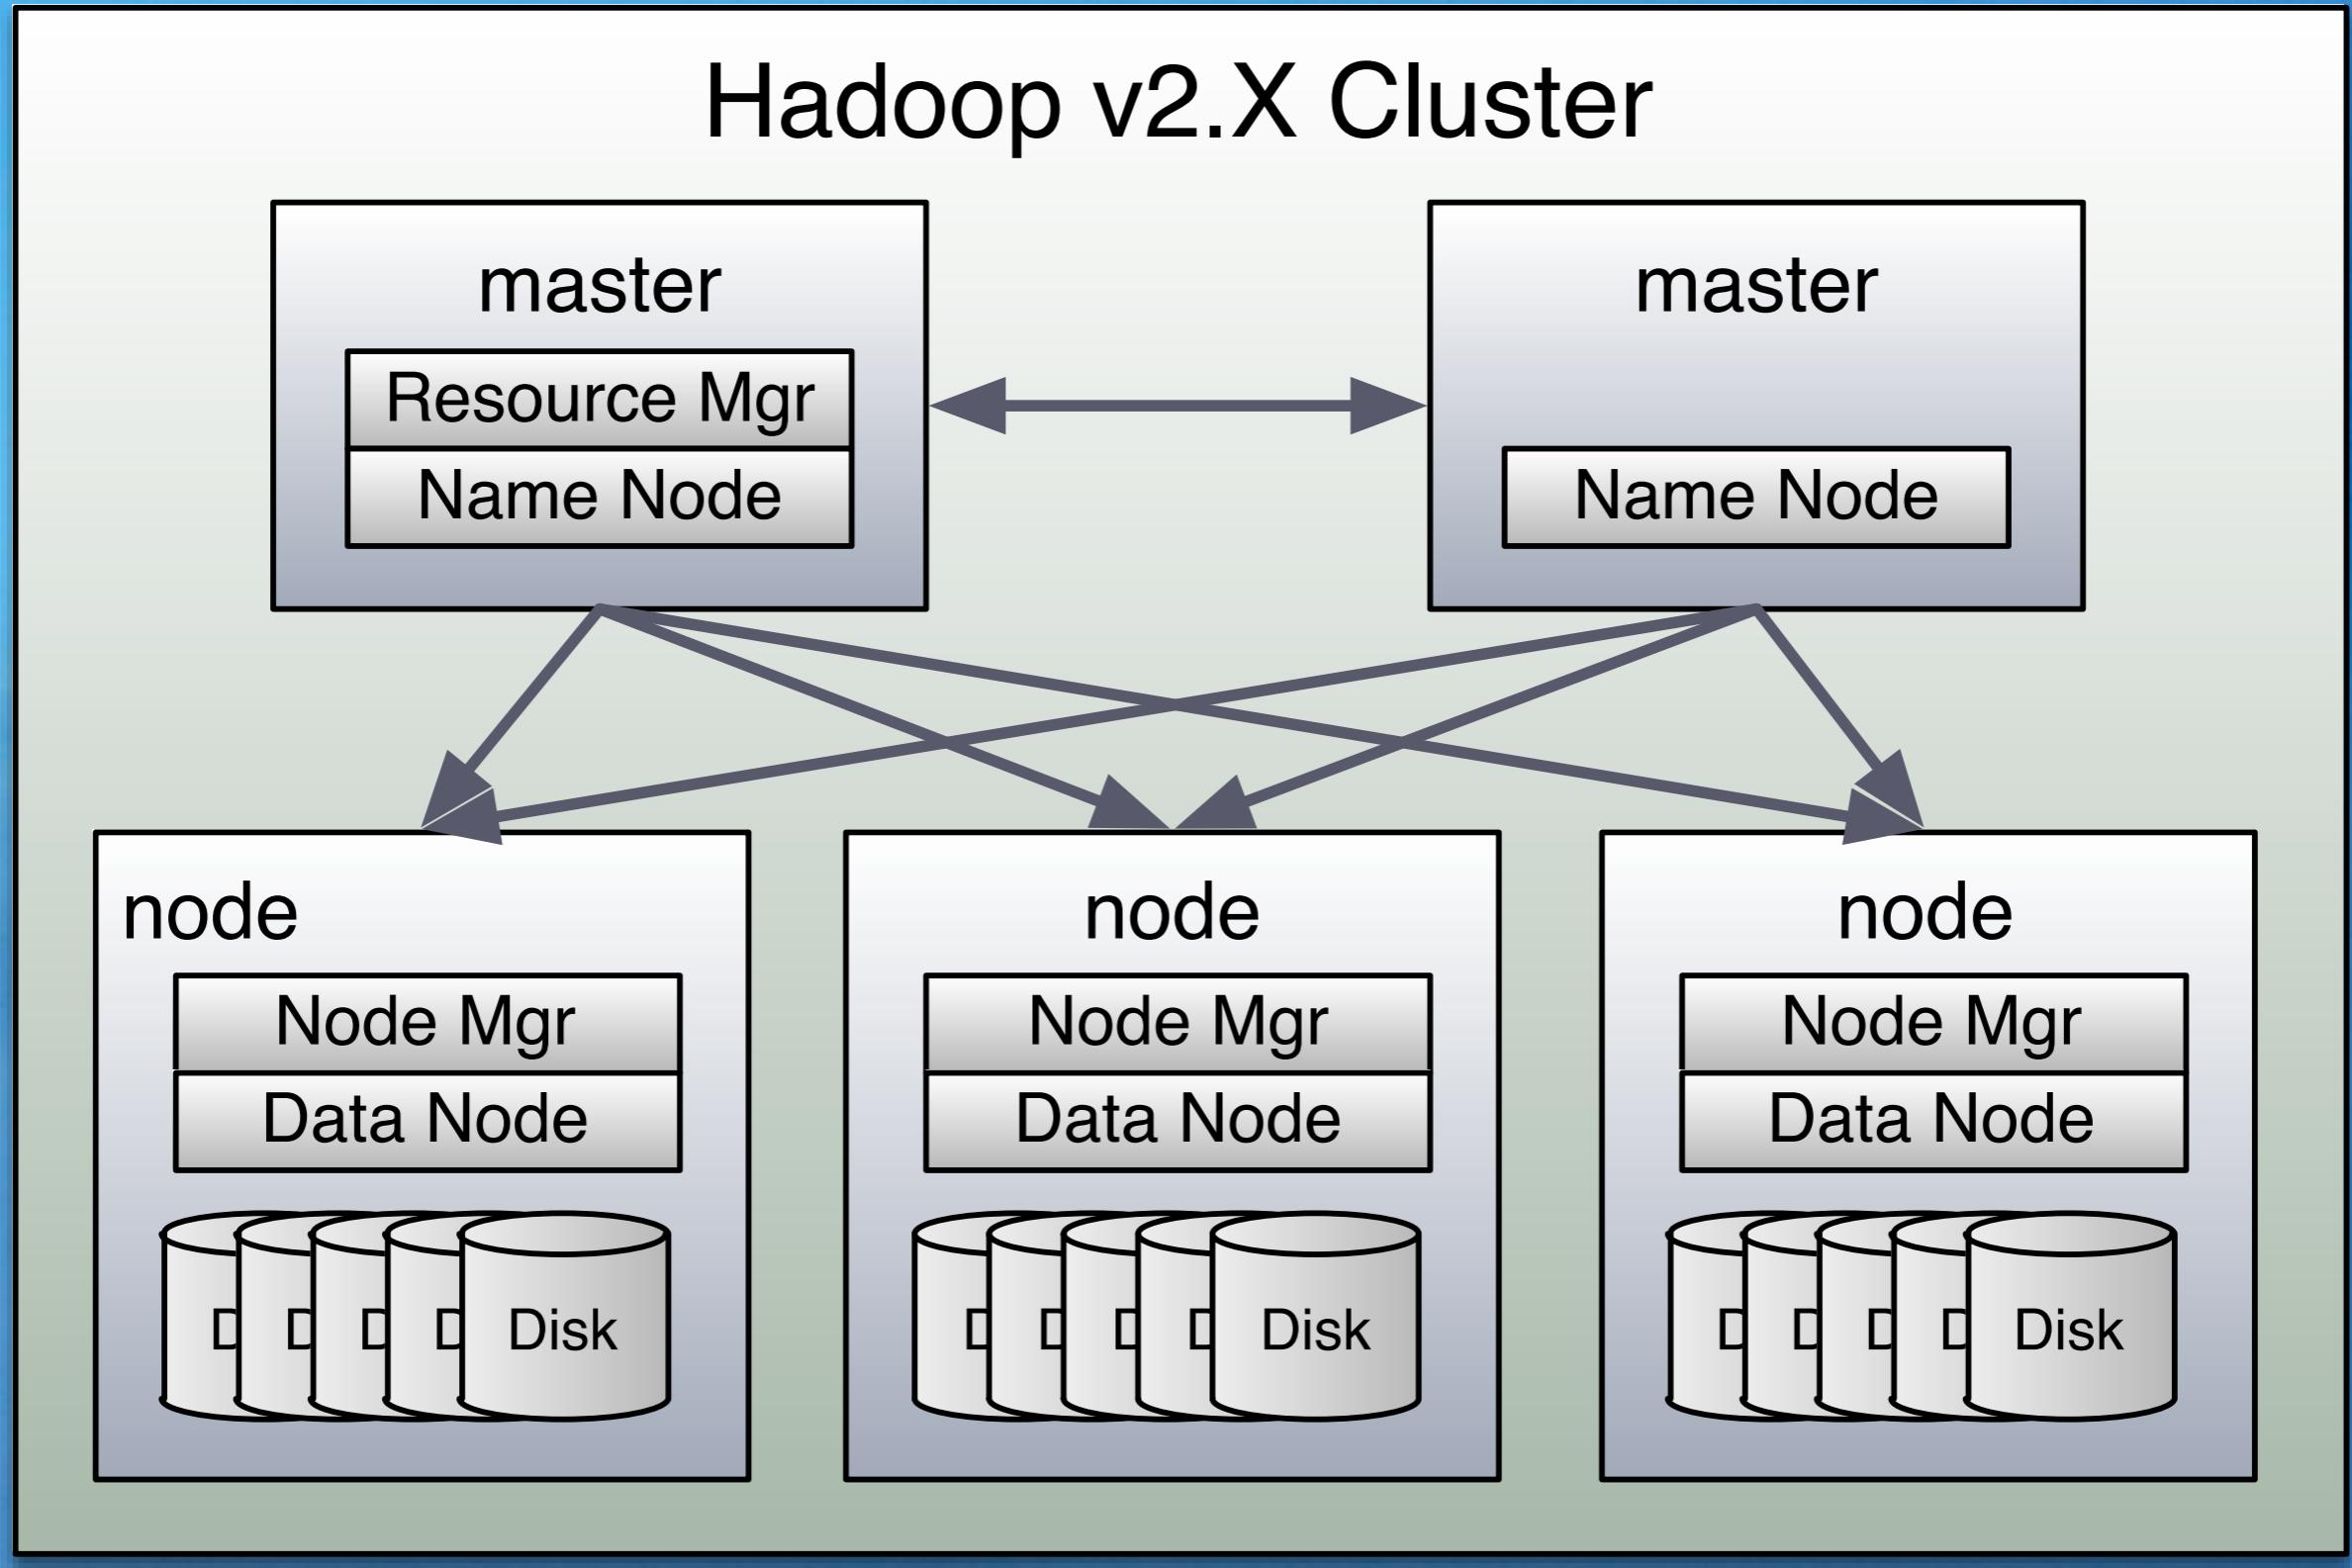


Wednesday, April 23, 14

The state of Hadoop as of last year.

Image: Detail of the London Eye

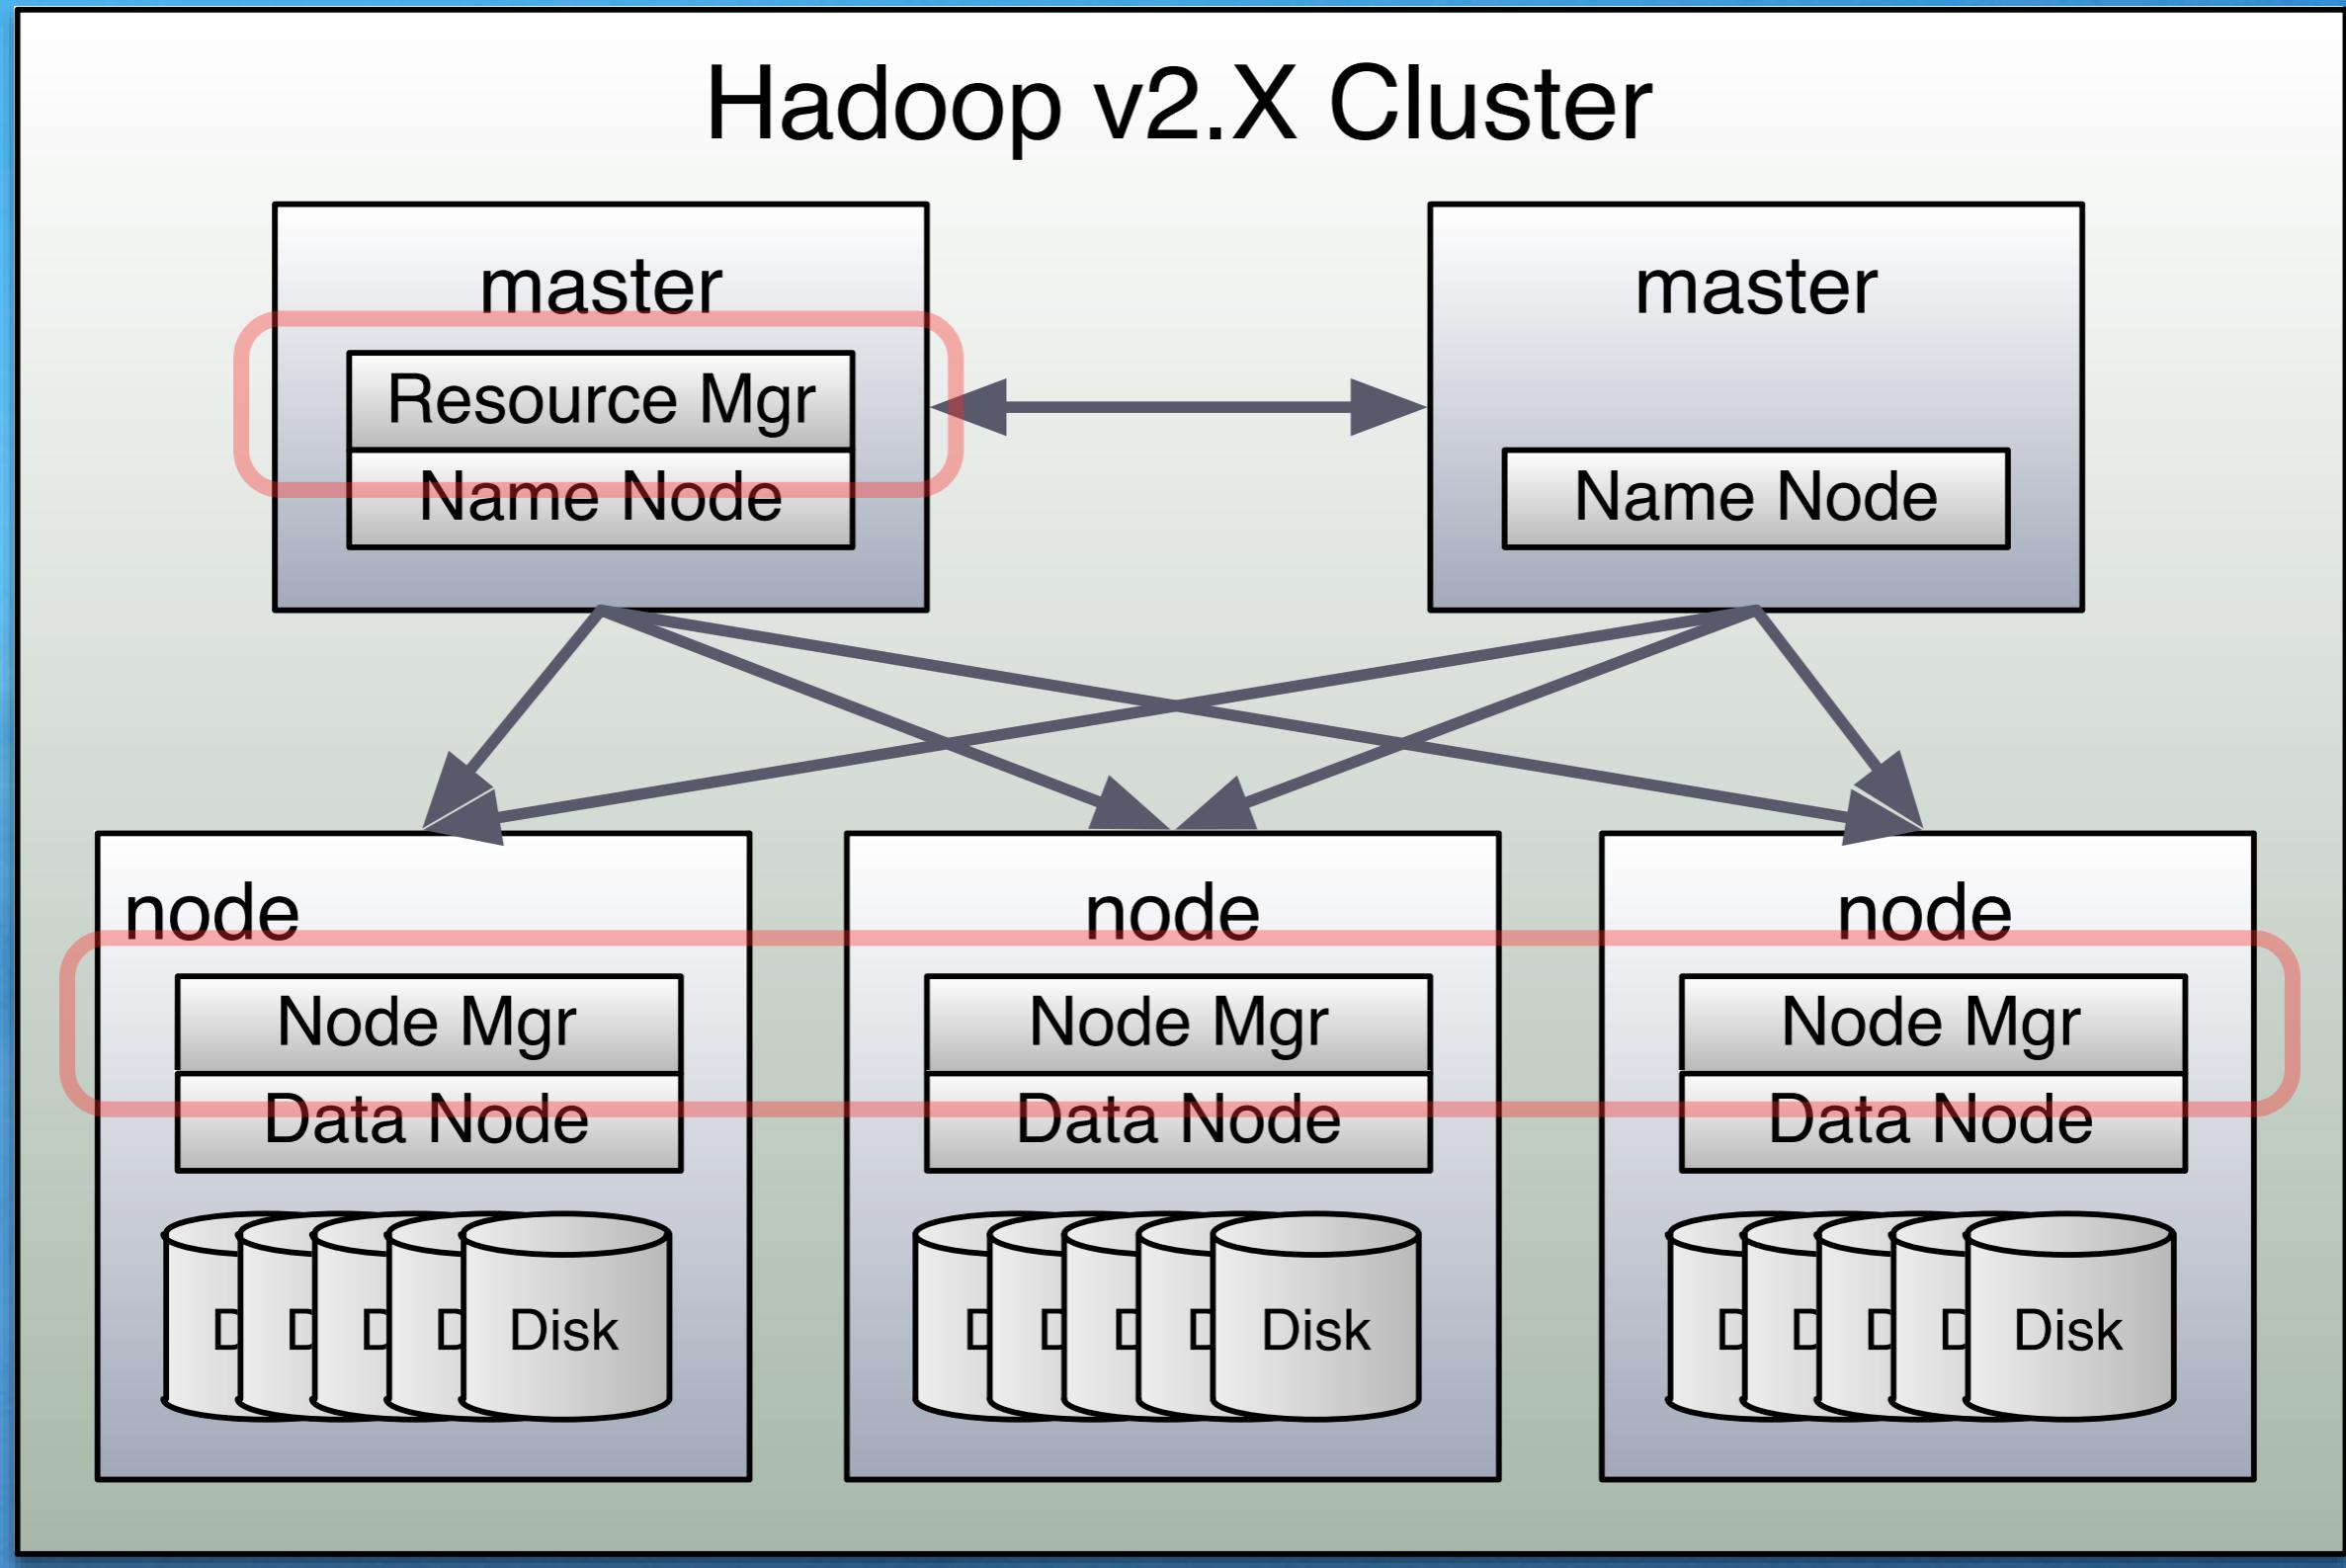
Hadoop v2.X Cluster



Wednesday, April 23, 14

Schematic view of a Hadoop 2 cluster. For a more precise definition of the services and what they do, see e.g., <http://hadoop.apache.org/docs/r2.3.0/hadoop-yarn/hadoop-yarn-site/YARN.html> We aren't interested in great details at this point, but we'll call out a few useful things to know.

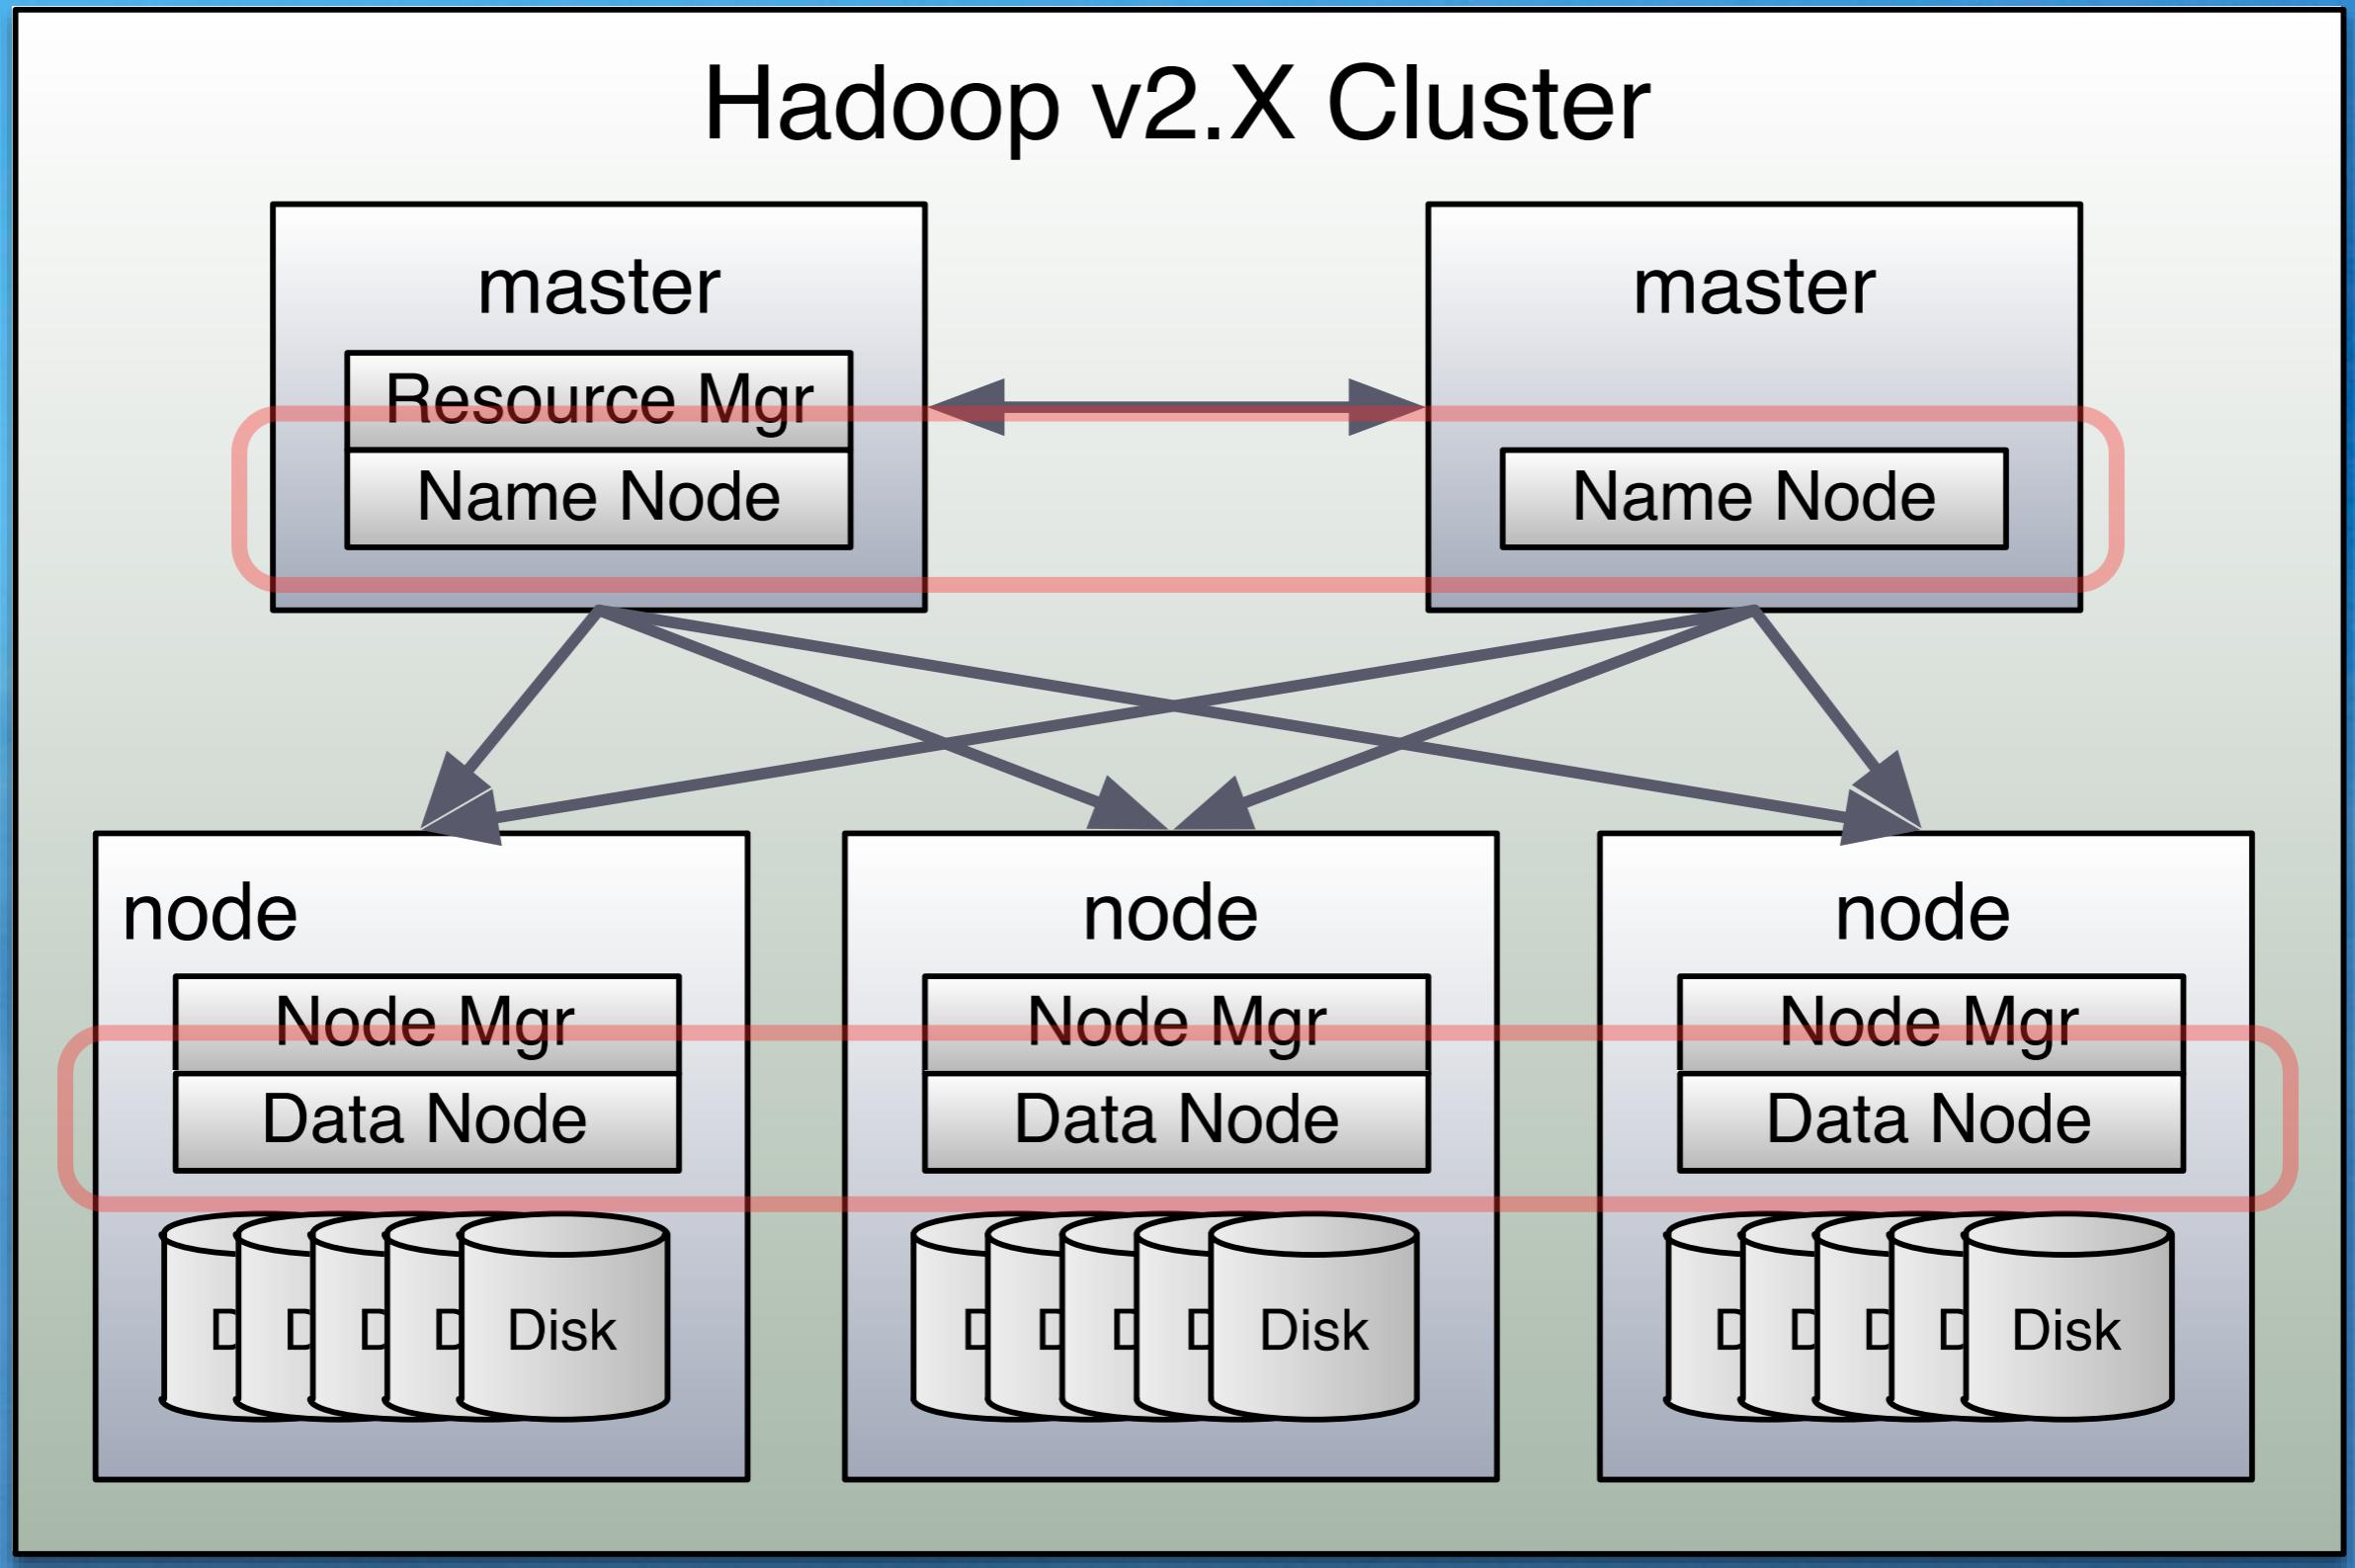
Resource and Node Managers



Wednesday, April 23, 14

Hadoop 2 uses YARN to manage resources via the master Resource Manager, which includes a pluggable job scheduler and an Applications Manager. It coordinates with the Node Manager on each node to schedule jobs and provide resources. Other services involved, including application-specific Containers and Application Masters are not shown.

Name Node and Data Nodes



Wednesday, April 23, 14

Hadoop 2 clusters federate the Name node services that manage the file system, HDFS. They provide horizontal scalability of file-system operations and resiliency when service instances fail. The data node services manage individual blocks for files.

MapReduce

The classic compute model
for Hadoop

Wednesday, April 23, 14

Hadoop 2 clusters federate the Name node services that manage the file system, HDFS. They provide horizontal scalability of file-system operations and resiliency when service instances fail. The data node services manage individual blocks for files.

MapReduce

**1 map step + 1 reduce step
(wash, rinse, repeat)**

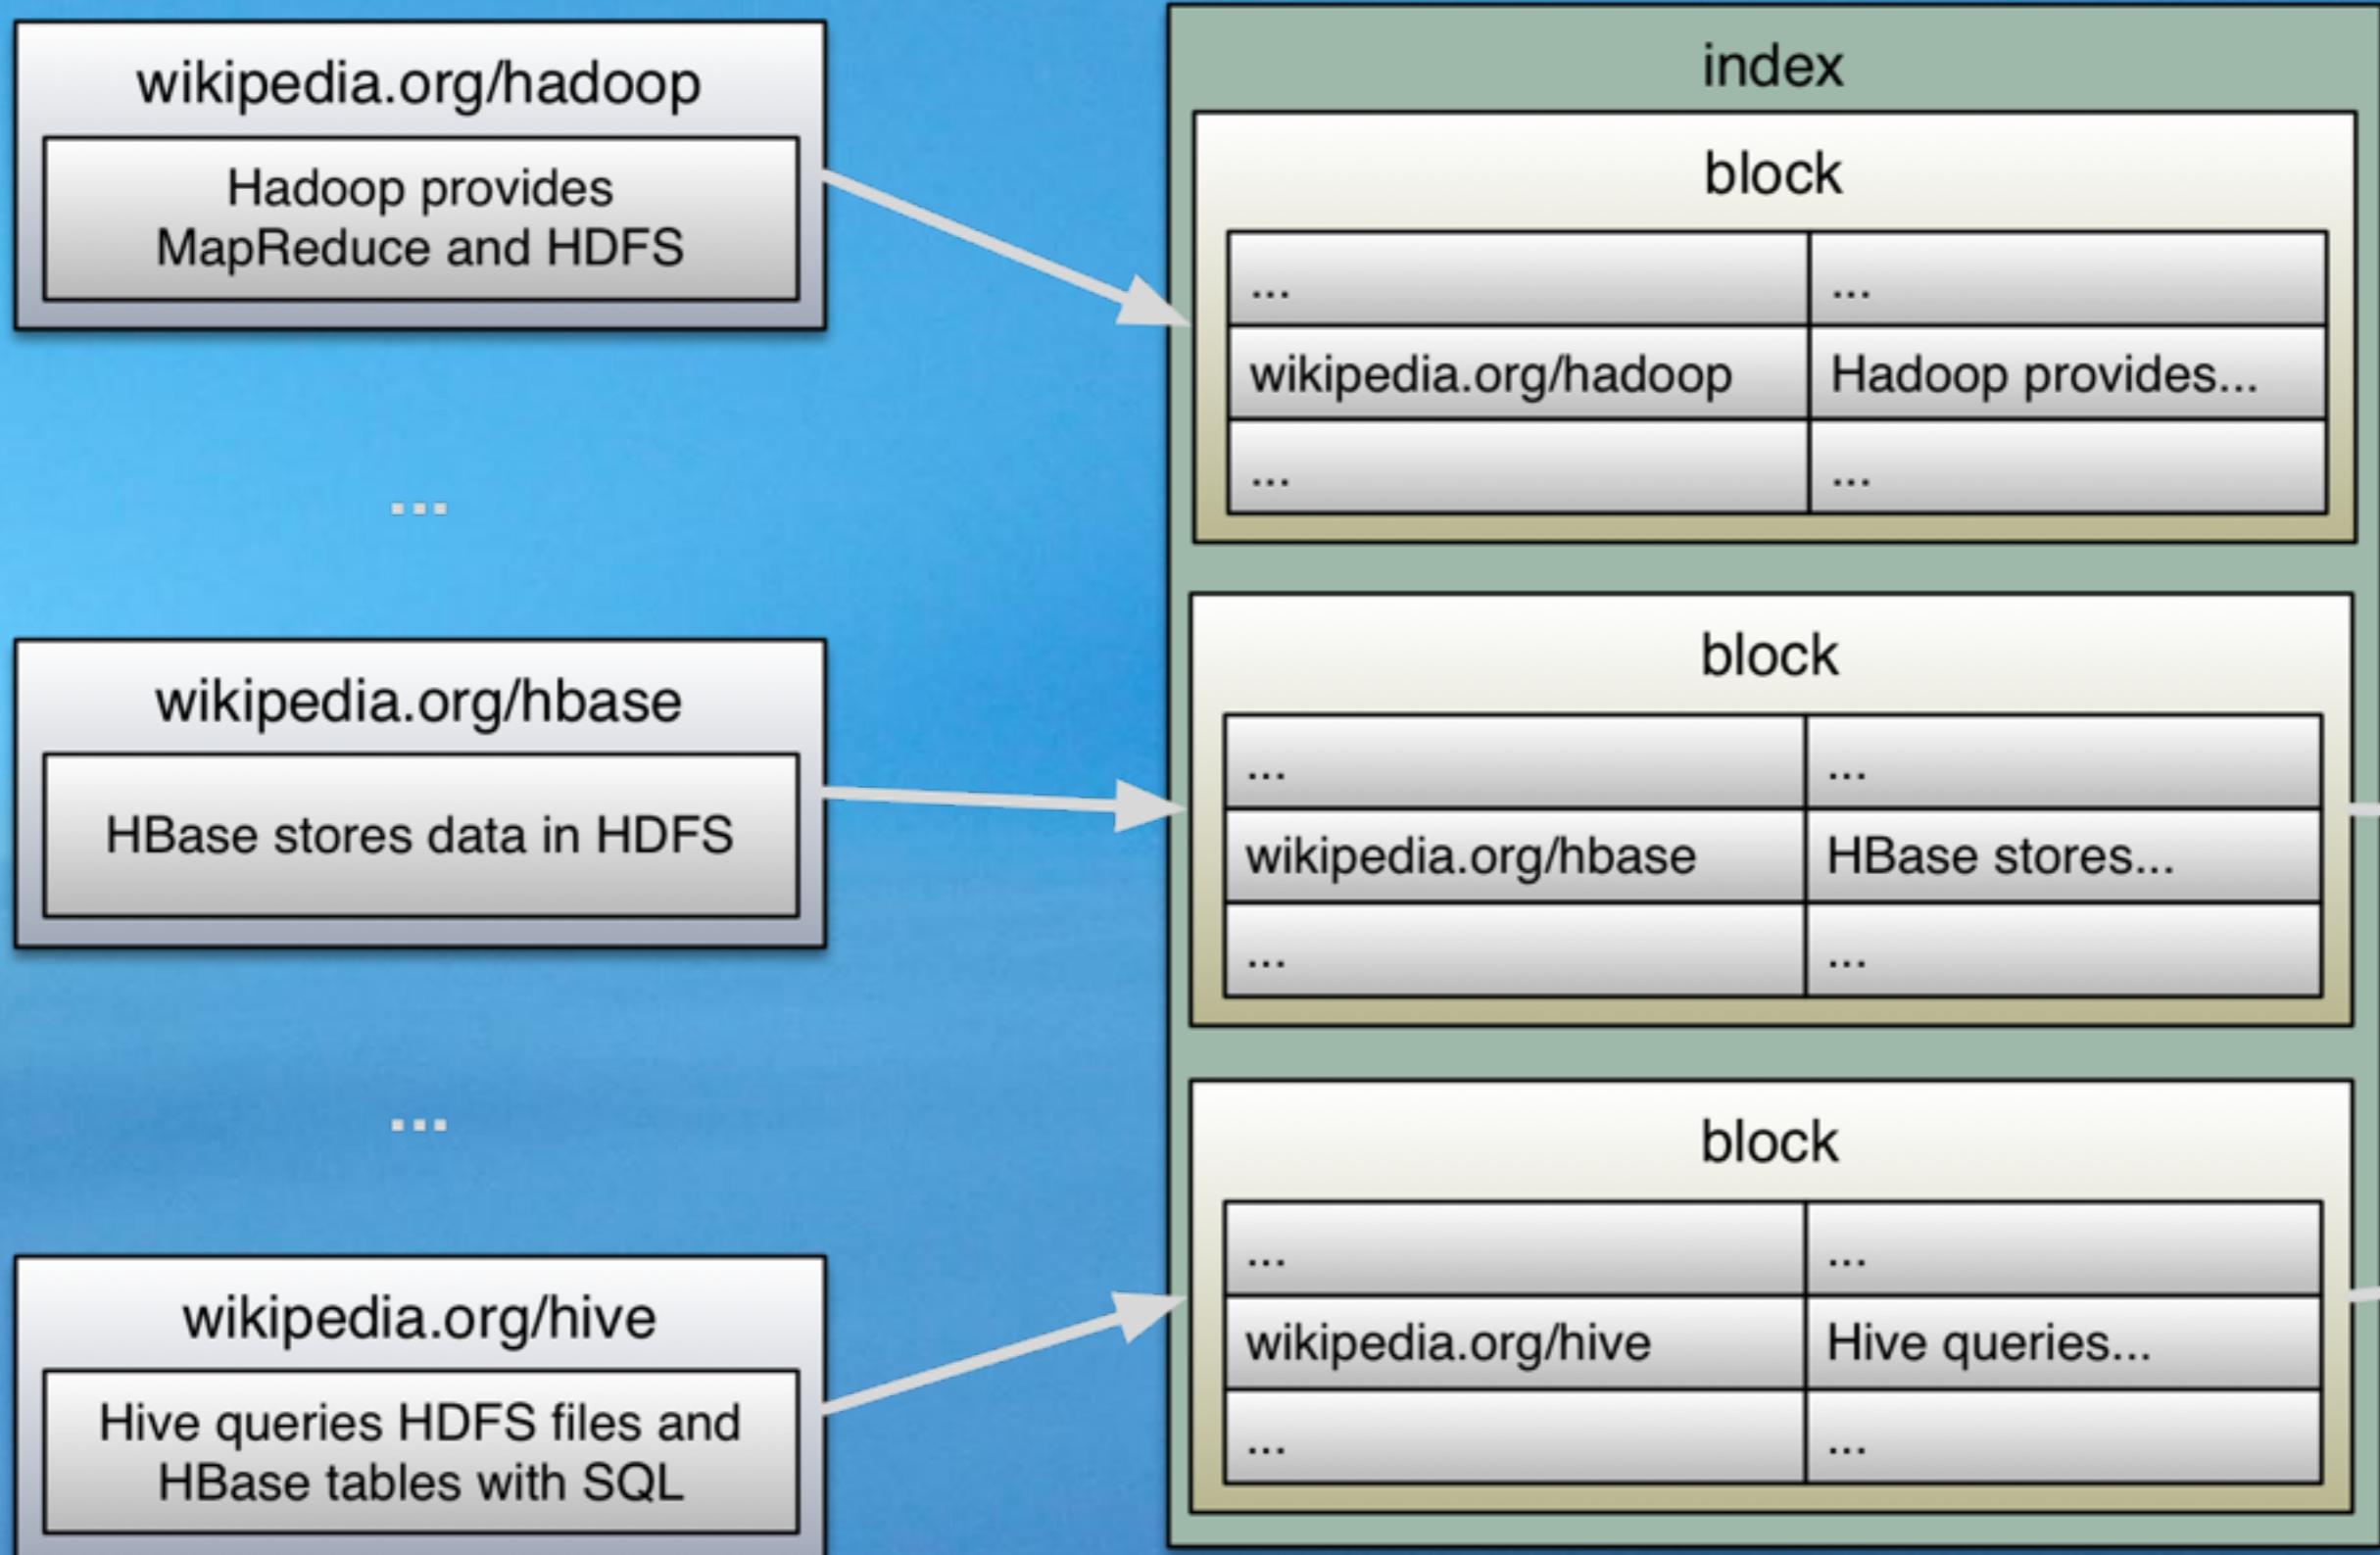
Wednesday, April 23, 14

You get 1 map step (although there is limited support for chaining mappers) and 1 reduce step. If you can't implement an algorithm in these two steps, you can chain jobs together, but you'll pay a tax of flushing the entire data set to disk between these jobs.

MapReduce

Example:
Inverted Index

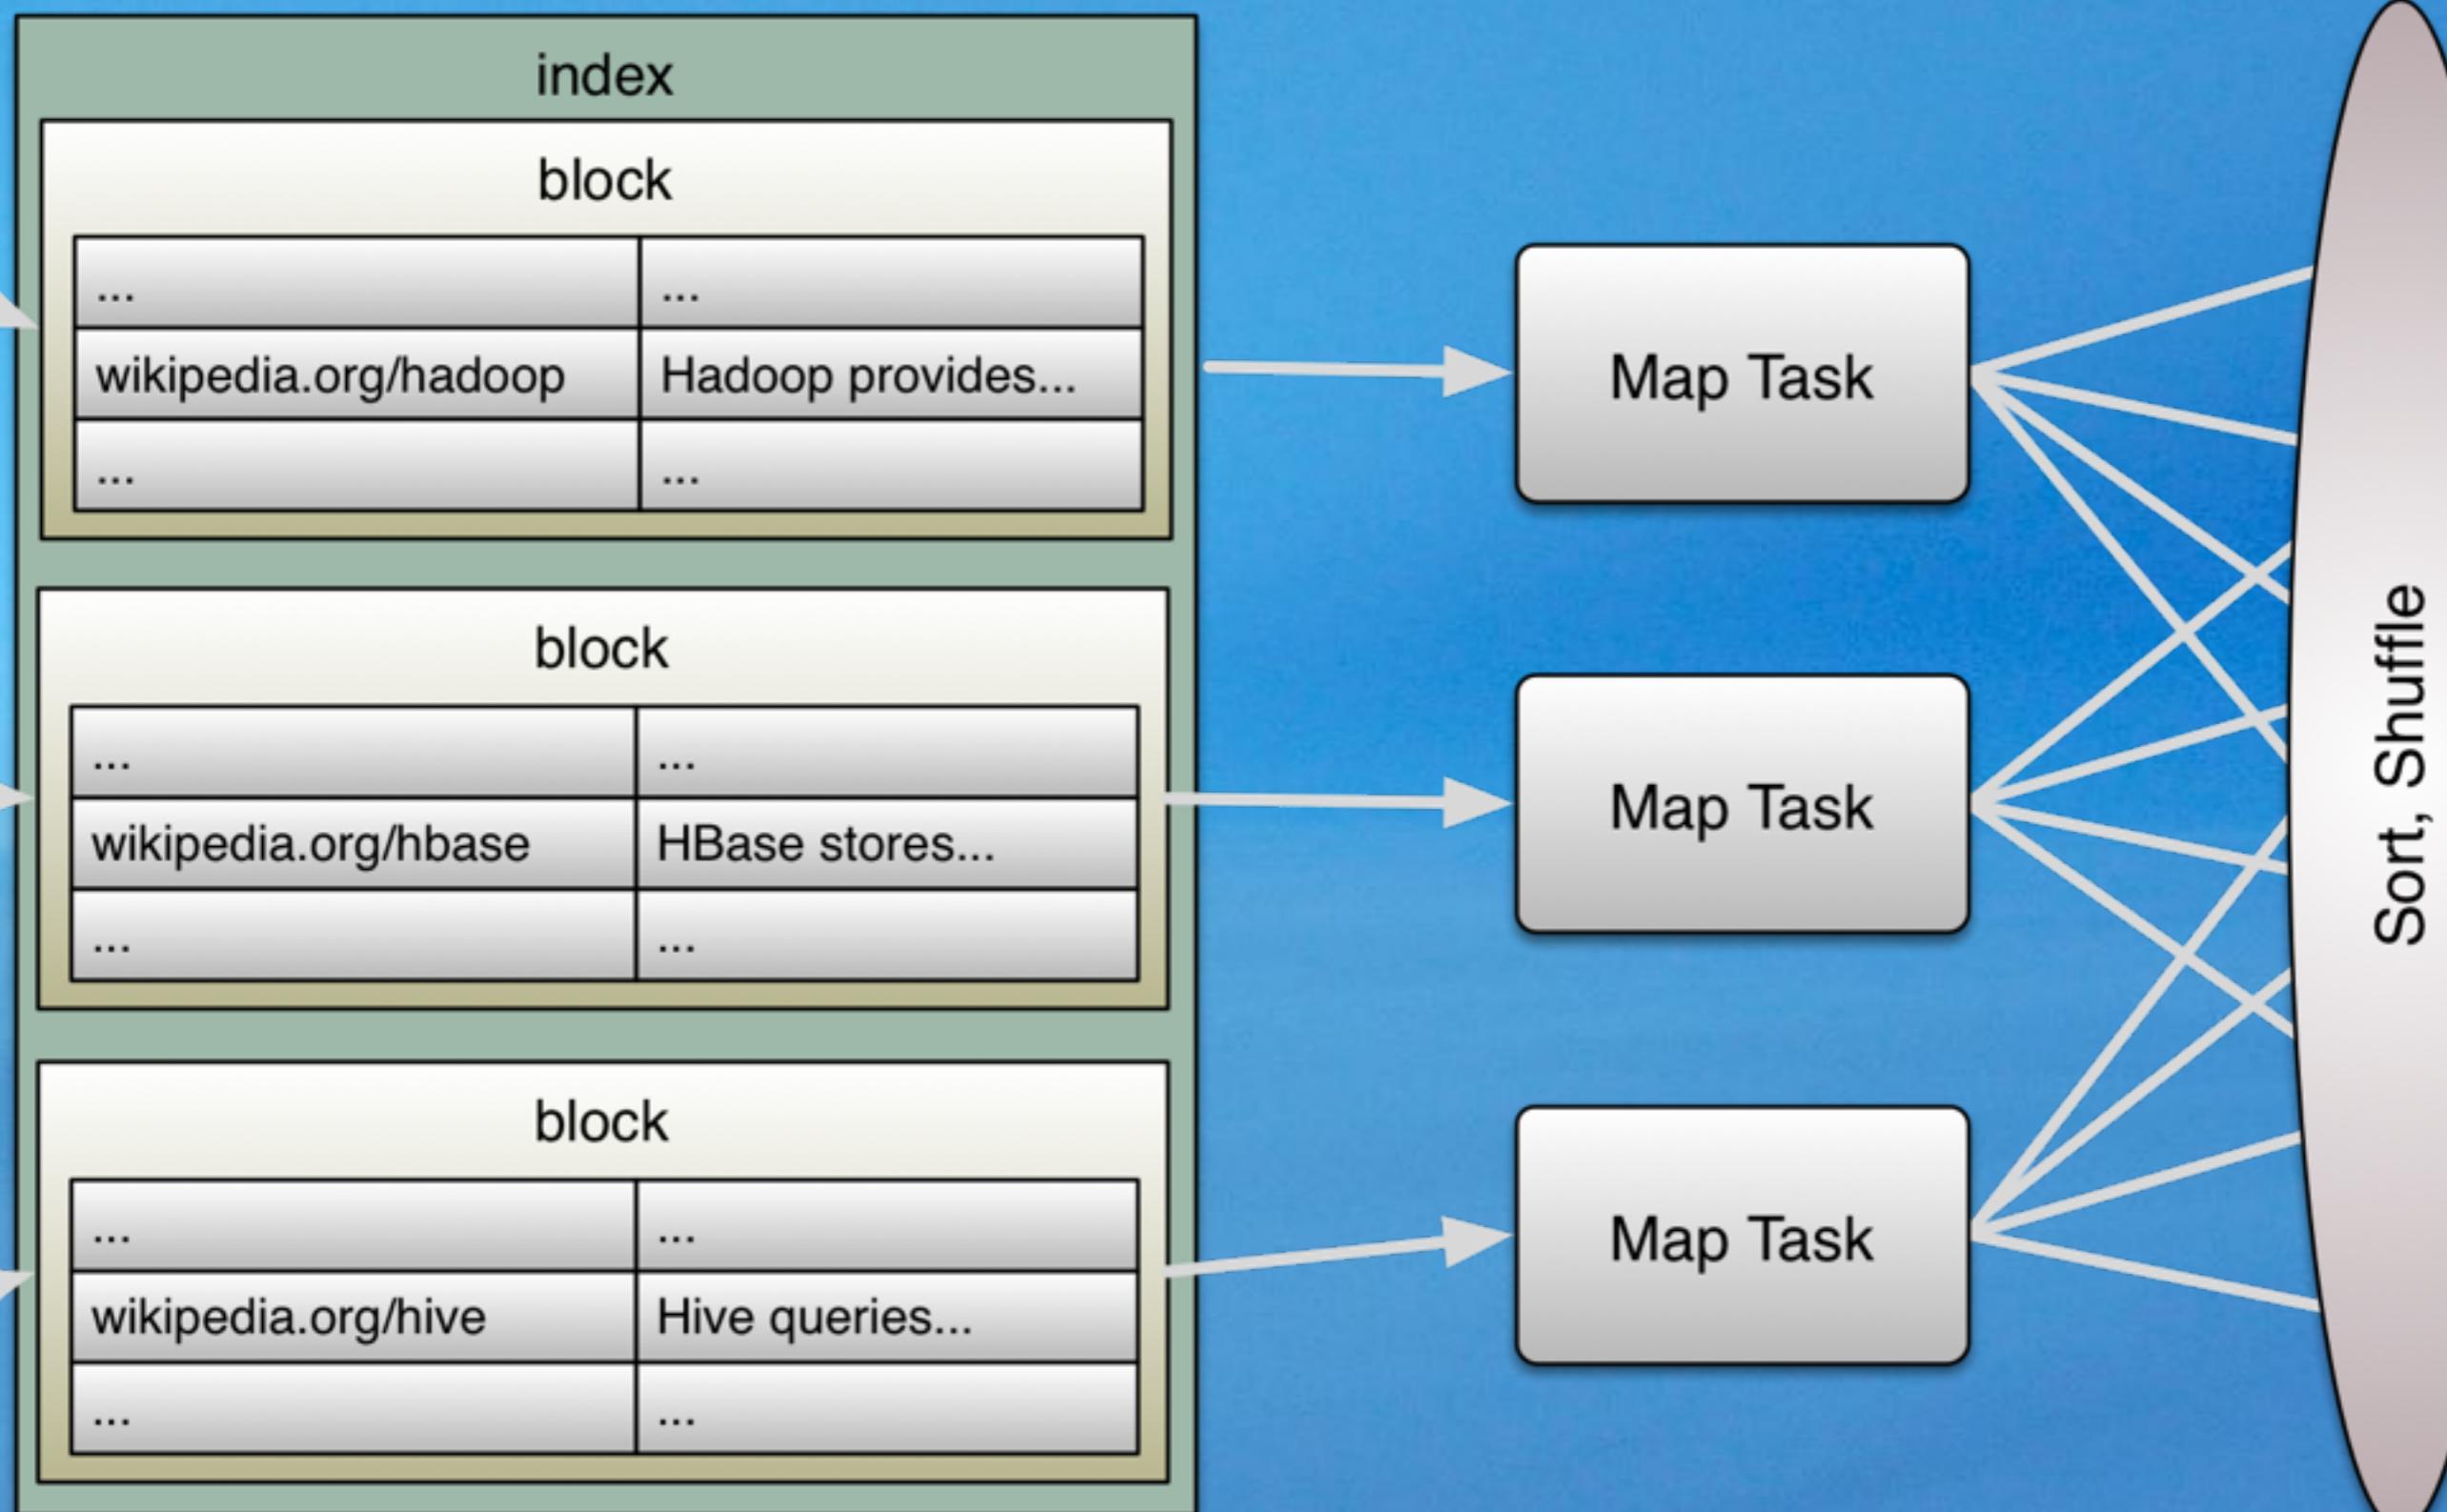
Web Crawl



Wednesday, April 23, 14

Before running MapReduce, crawl teh interwebs, find all the pages, and build a data set of URLs -> doc contents, written to flat files in HDFS or one of the more “sophisticated” formats.

Map Phase



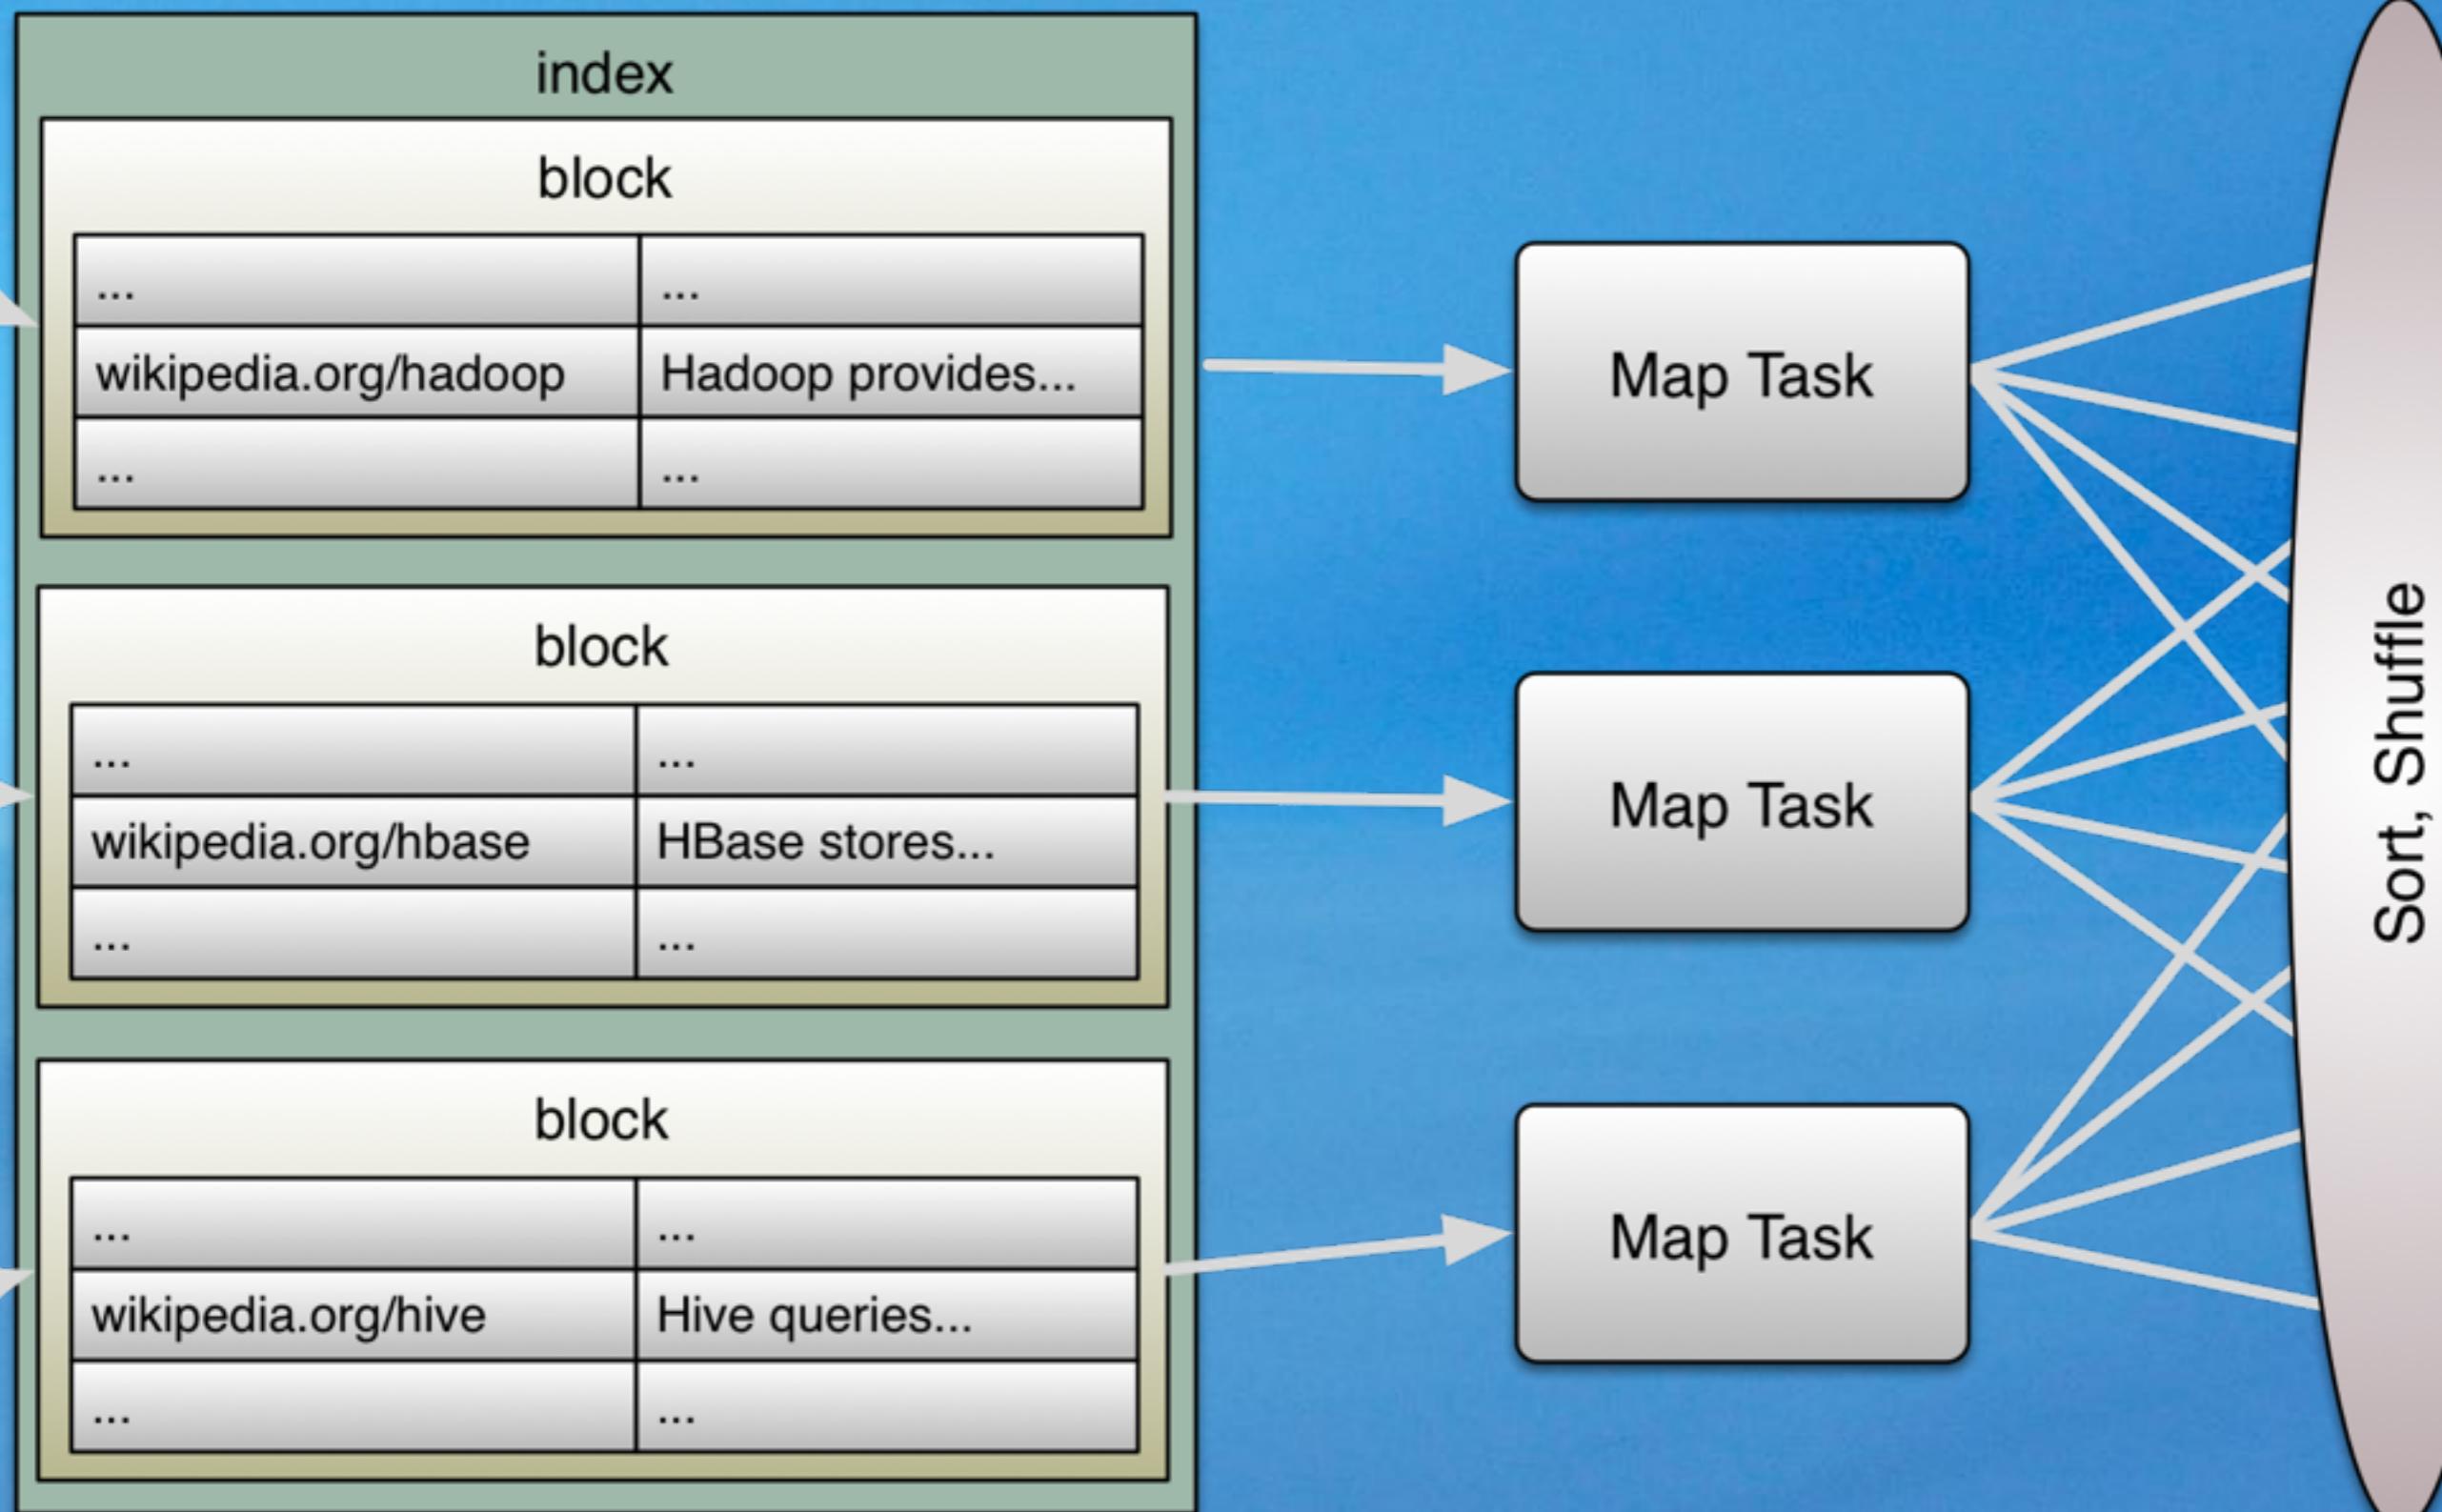
Wednesday, April 23, 14

Now we're running MapReduce. In the map step, a task (JVM process) per file *block* (64MB or larger) reads the rows, tokenizes the text and outputs key-value pairs ("tuples")...

Map Task

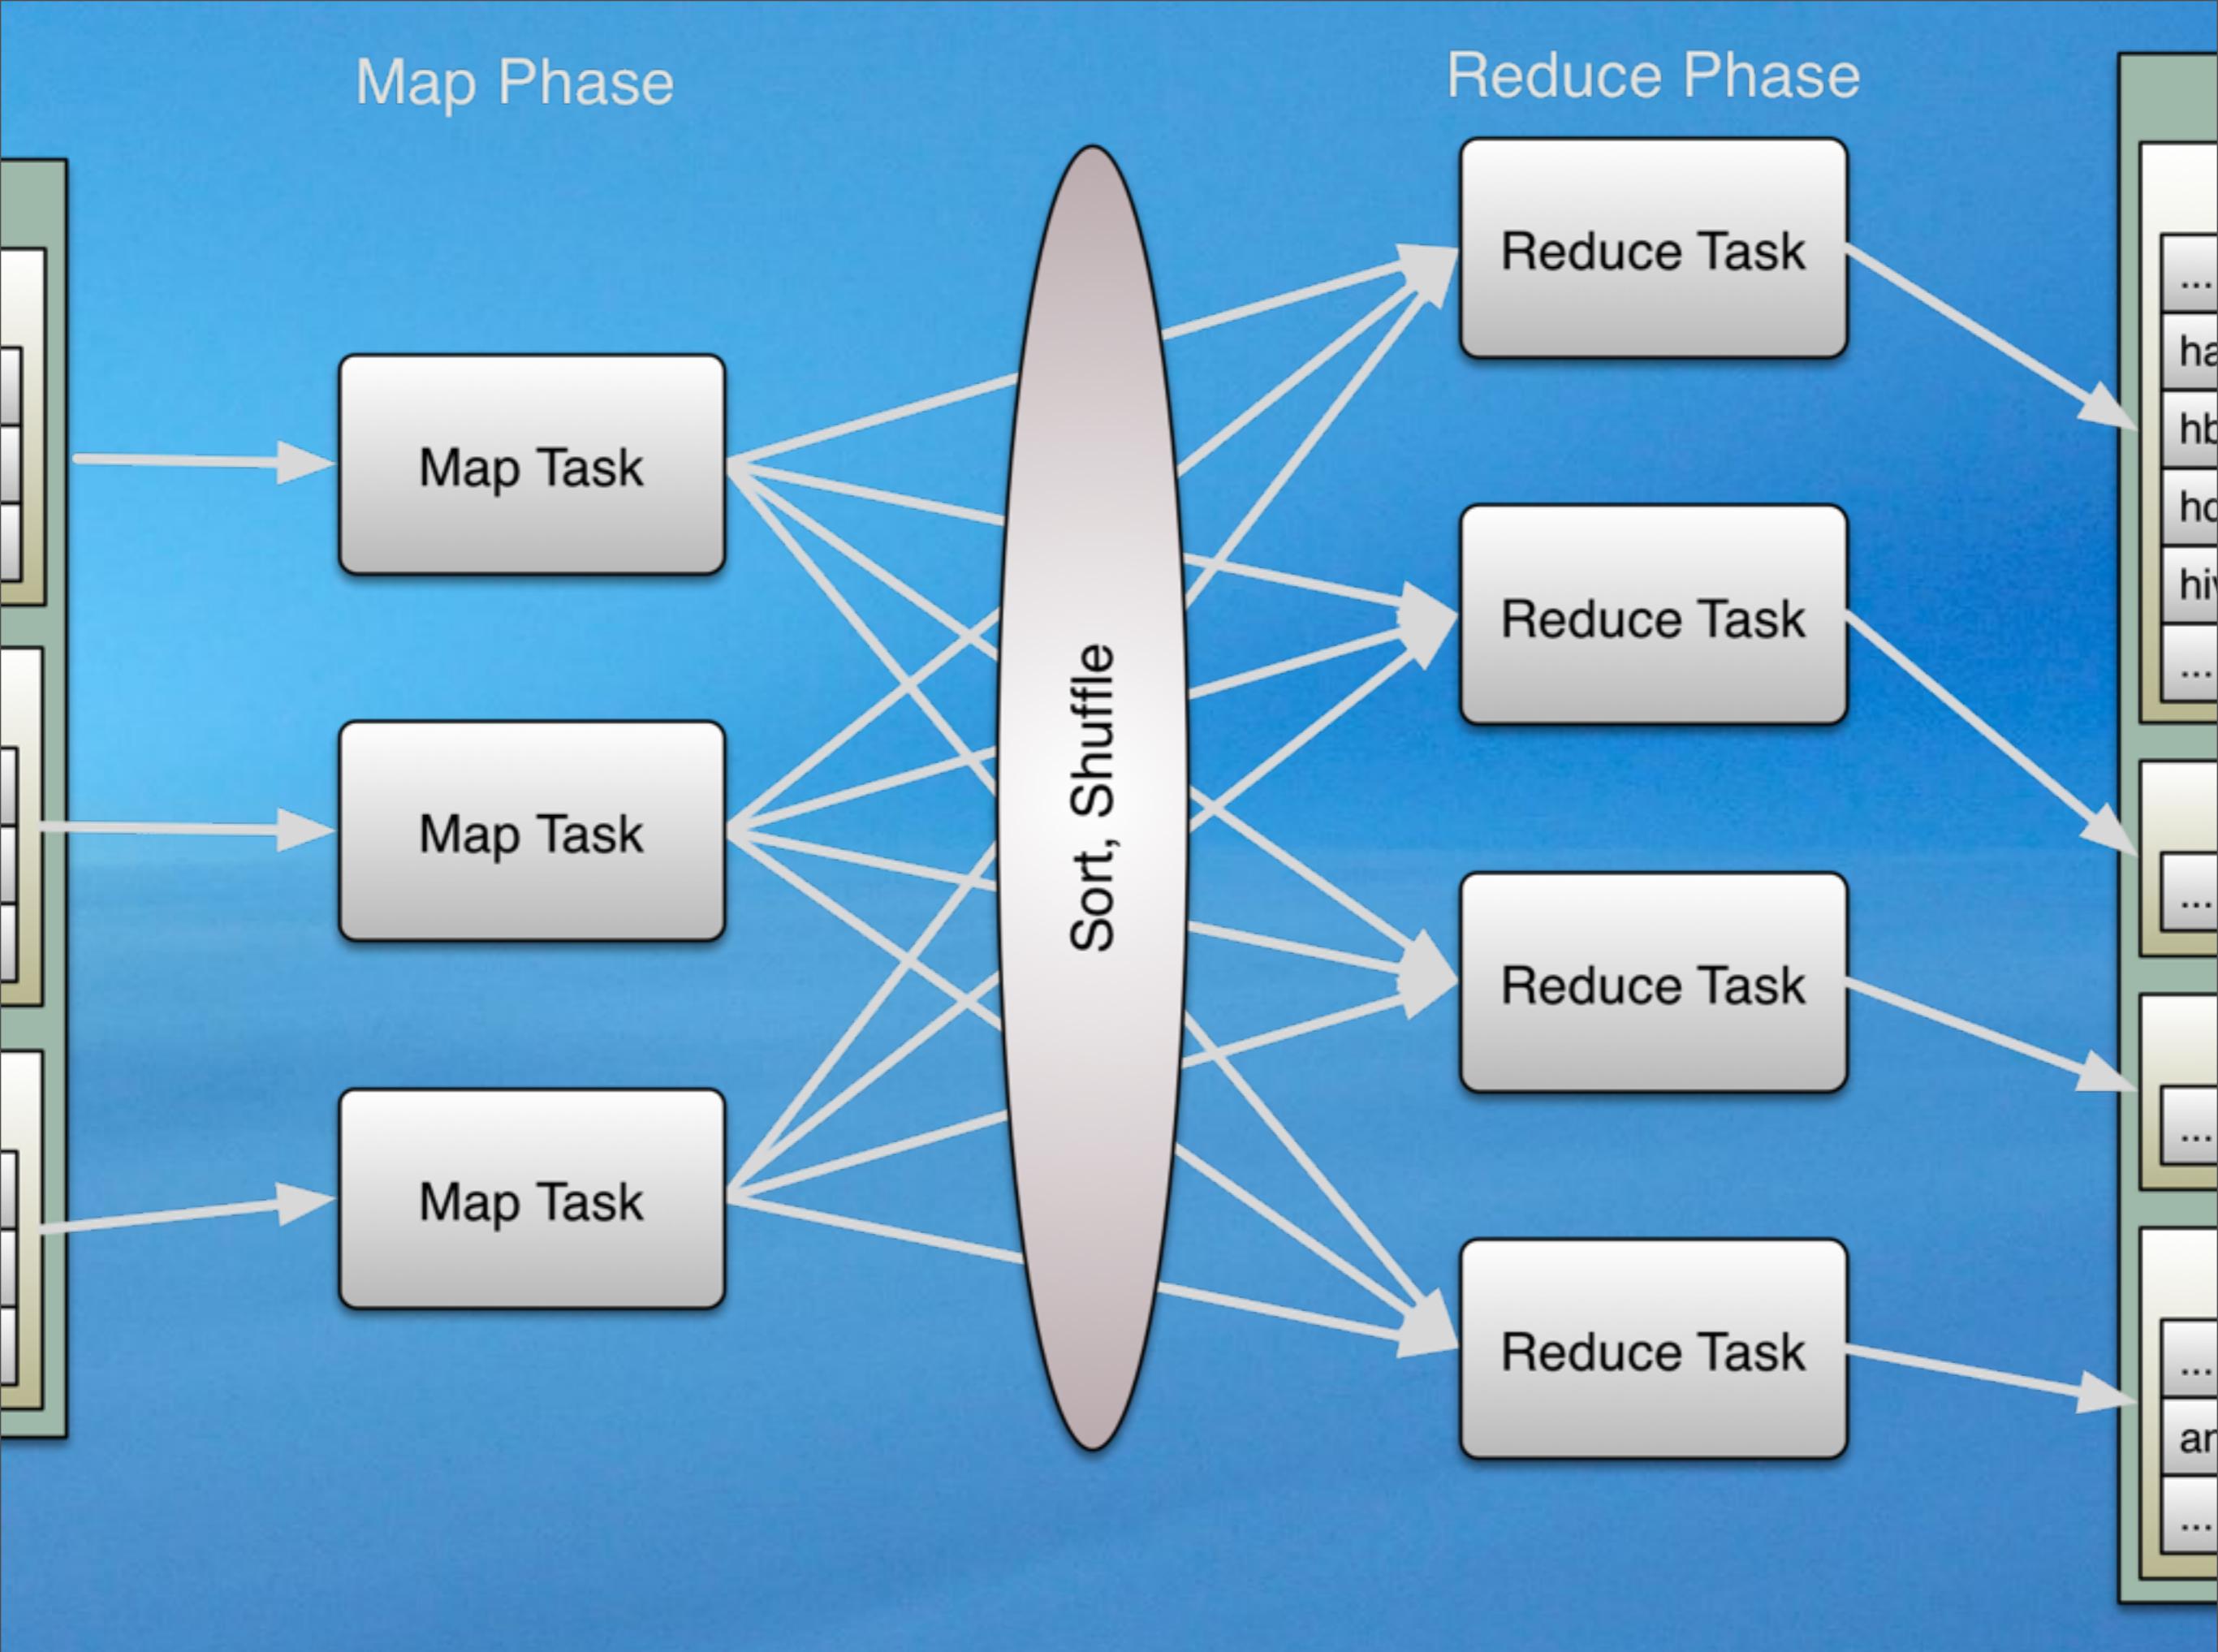
(hadoop,(wikipedia.org/hadoop,1))
(provides,(wikipedia.org/hadoop,1))
(mapreduce,(wikipedia.org/hadoop, 1))
(and,(wikipedia.org/hadoop,1))
(hdfs,(wikipedia.org/hadoop, 1))

Map Phase



Map Phase

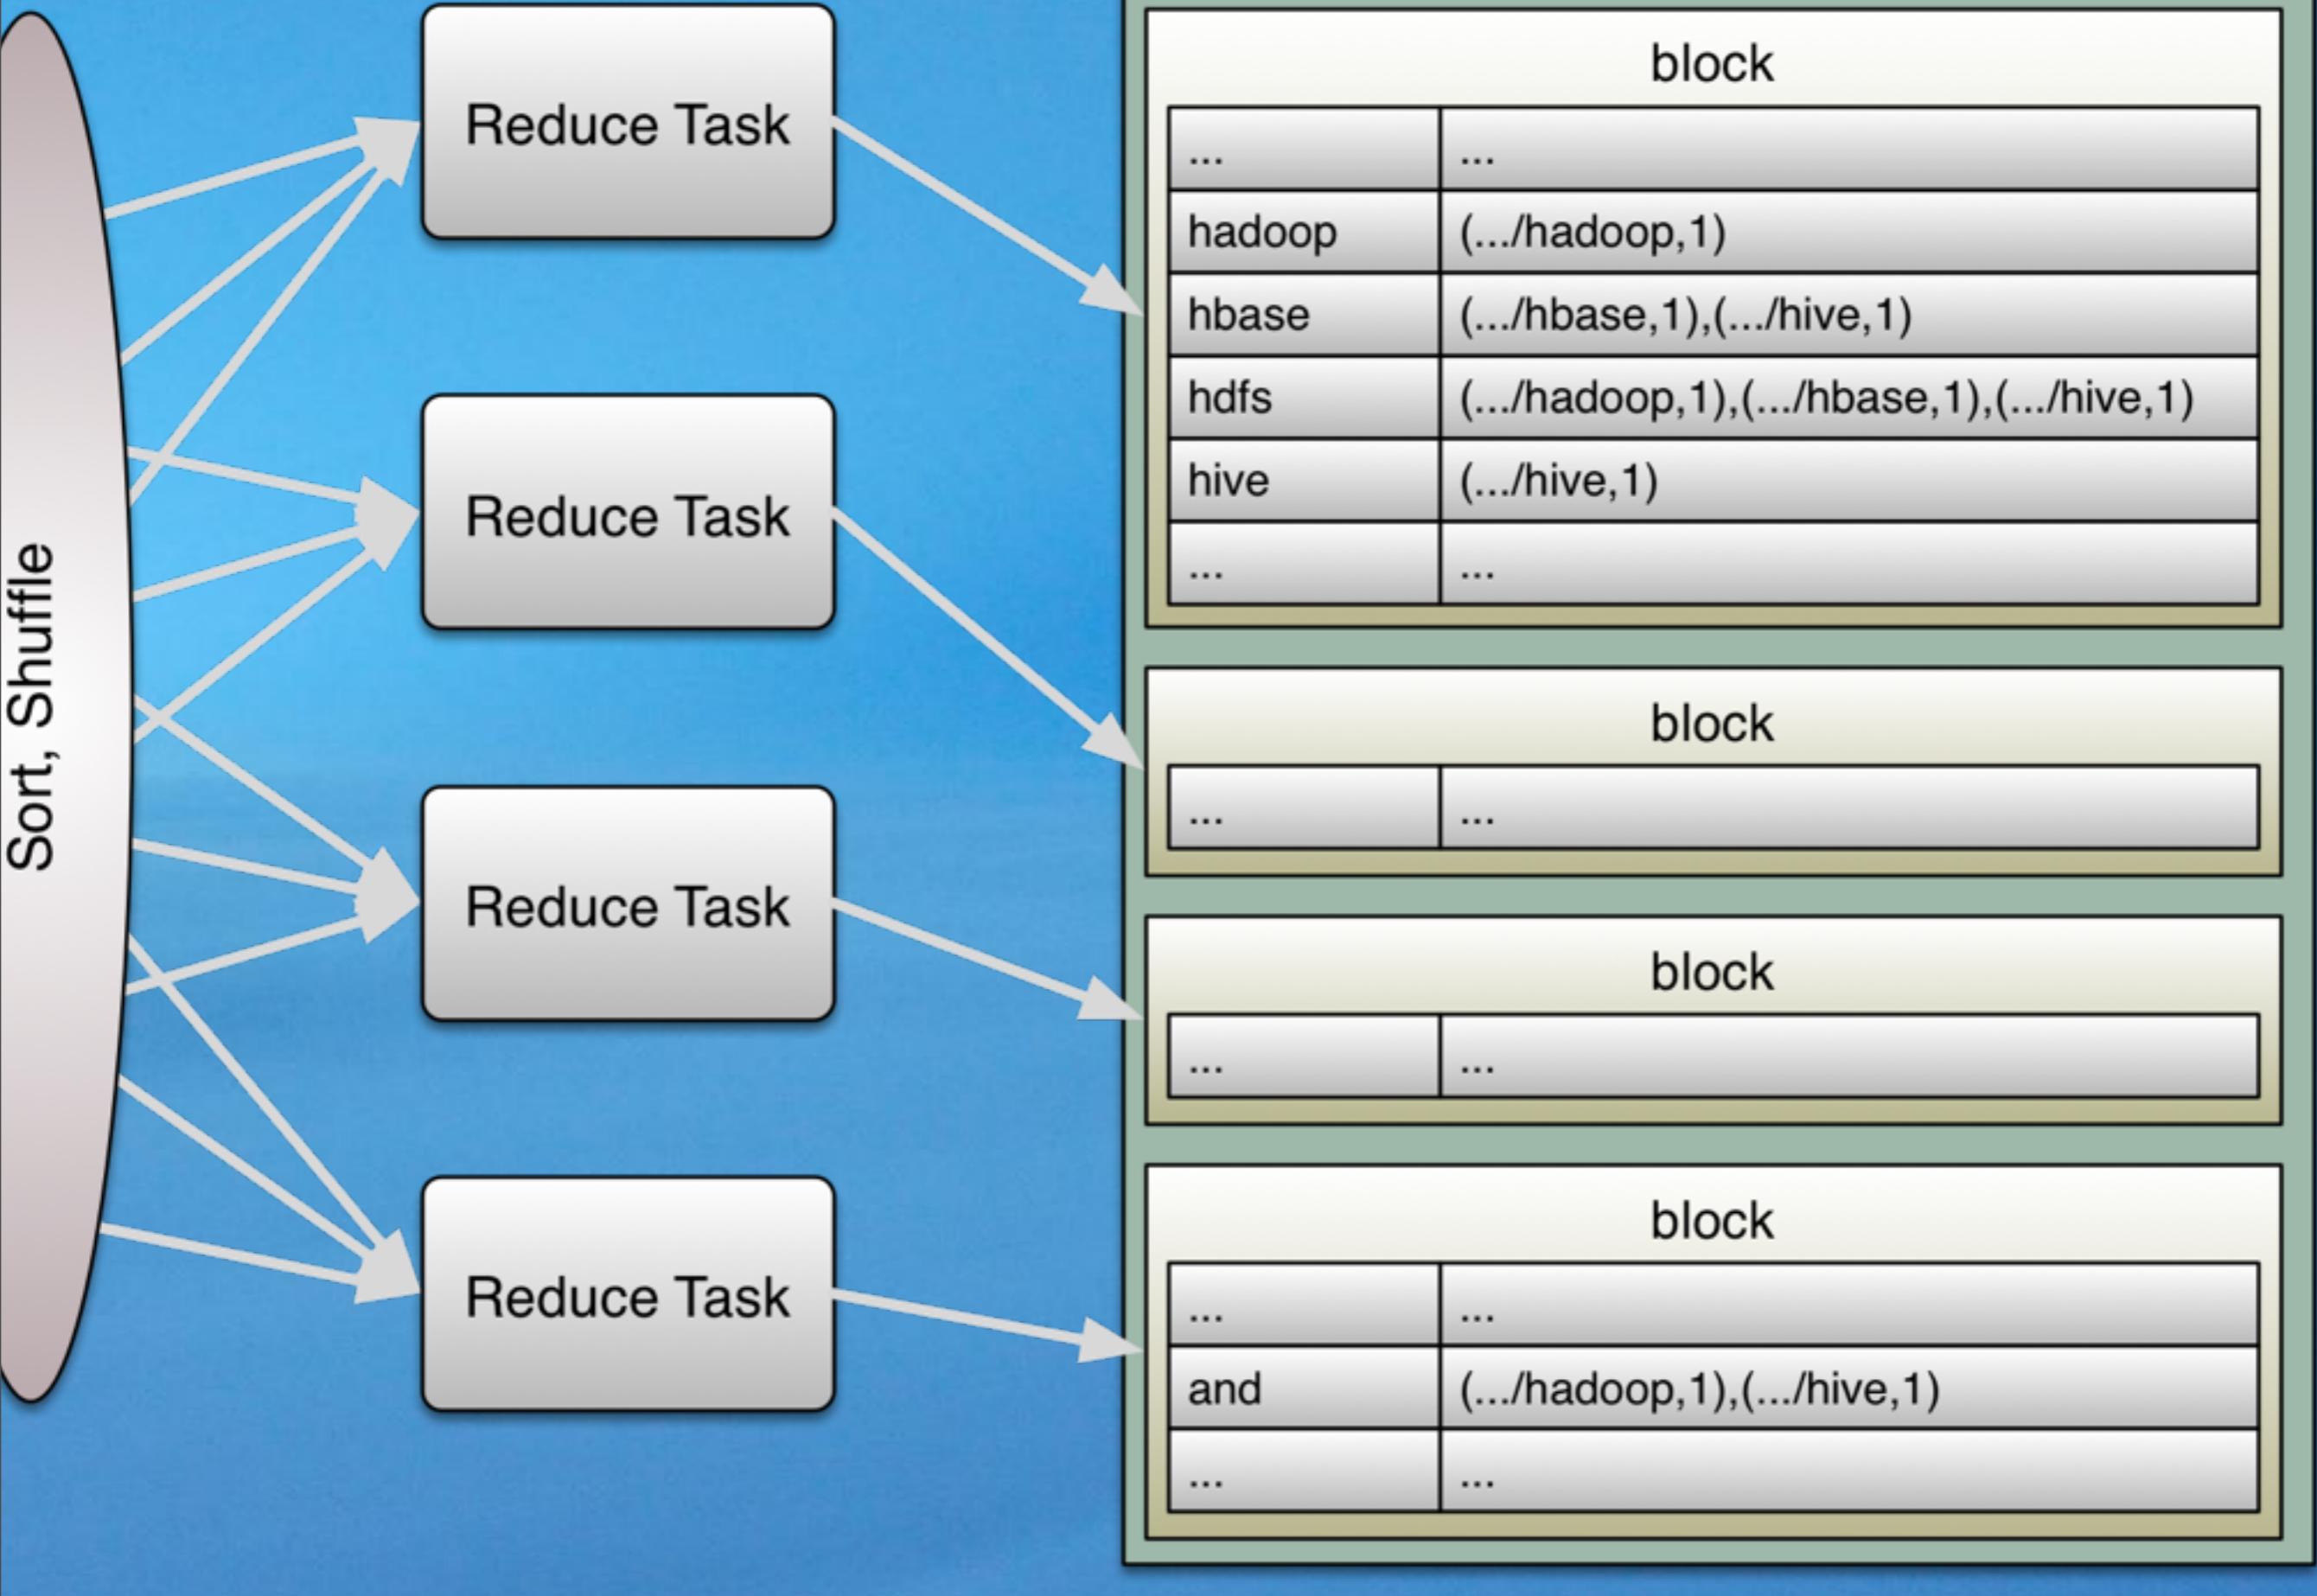
Reduce Phase



Wednesday, April 23, 14

The output tuples are sorted by key locally in each map task, then “shuffled” over the cluster network to reduce tasks (each a JVM process, too), where we want all occurrences of a given key to land on the same reduce task.

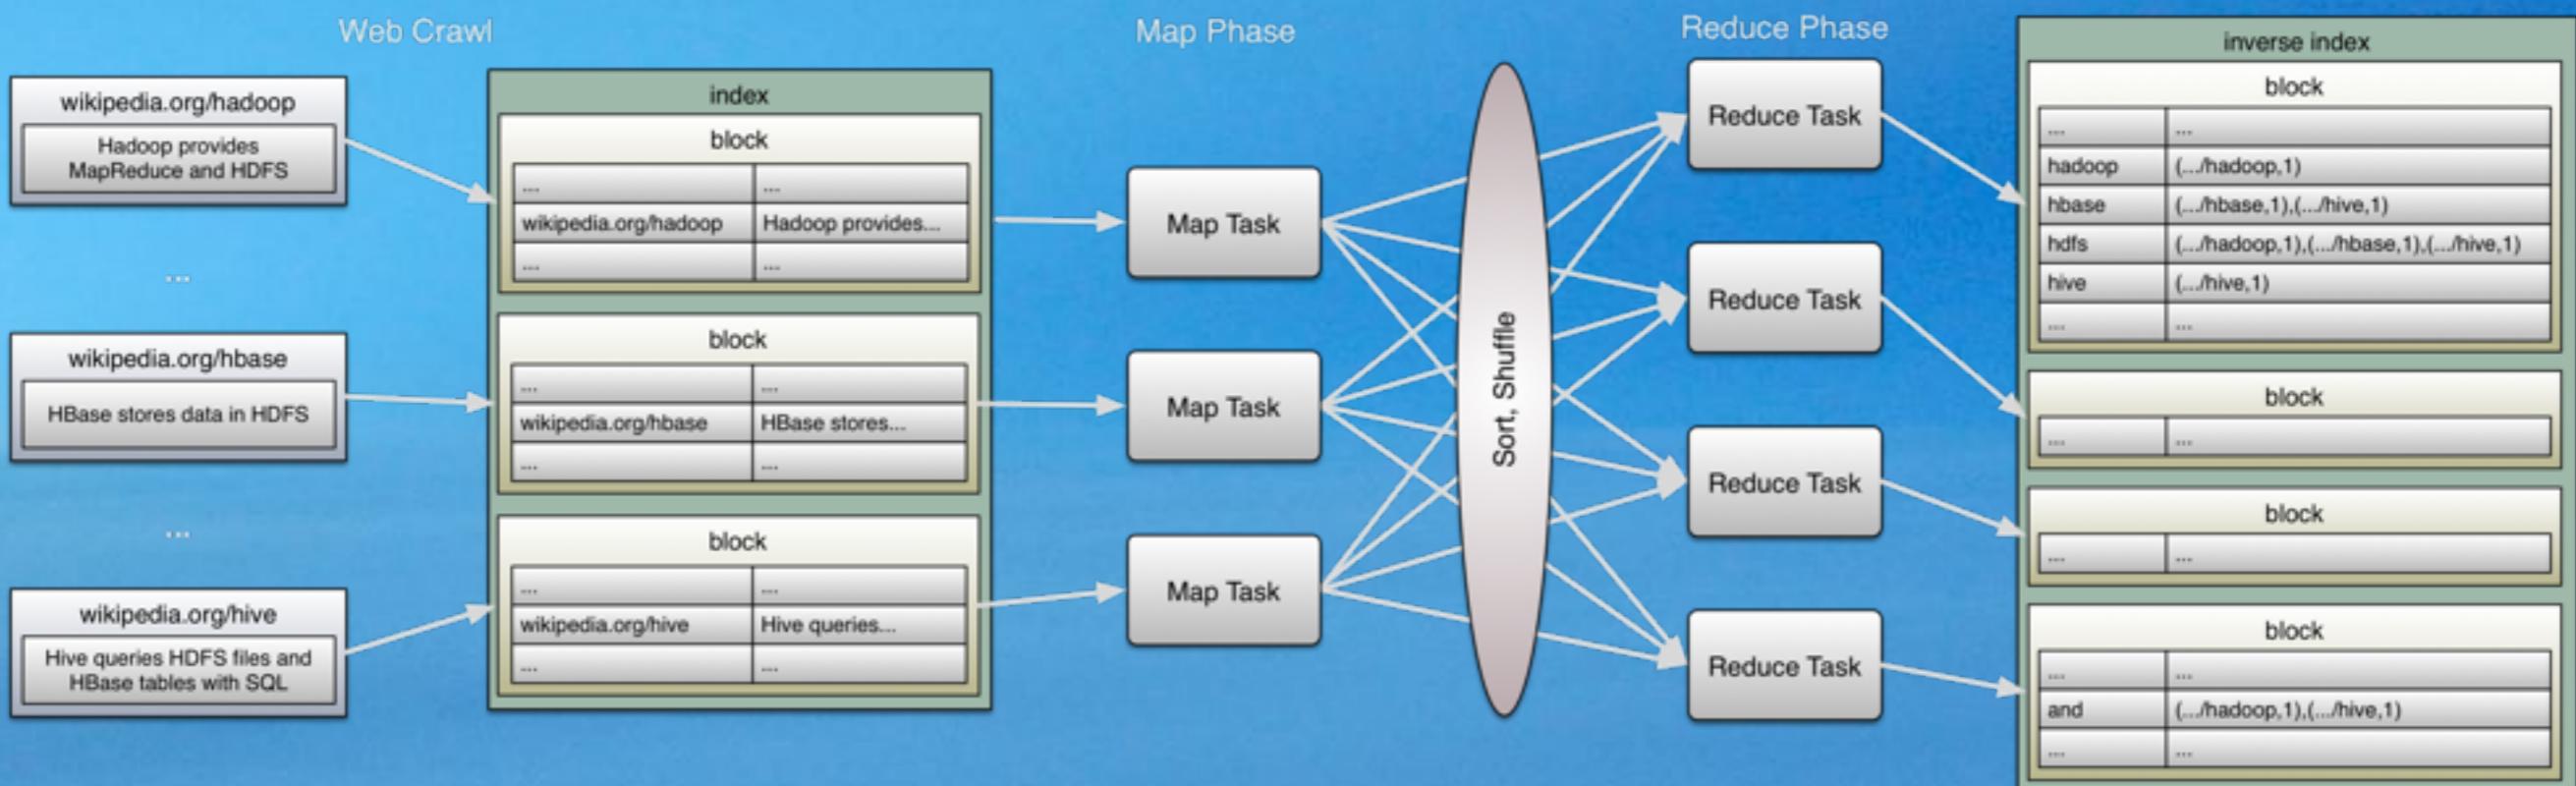
Reduce Phase



Wednesday, April 23, 14

Finally, each reducer just aggregates all the values it receives for each key, then writes out new files to HDFS with the words and a list of (URL-count) tuples (pairs).

Altogether...



Wednesday, April 23, 14

Finally, each reducer just aggregates all the values it receives for each key, then writes out new files to HDFS with the words and a list of (URL-count) tuples (pairs).



What's
not to like?

Wednesday, April 23, 14

This seems okay, right? What's wrong with it?

Awkward

Most algorithms are
much harder to implement
in this restrictive
map-then-reduce model.

Awkward

Lack of flexibility inhibits optimizations, too.

Wednesday, April 23, 14

The inflexible compute model leads to complex code to improve performance where hacks are used to work around the limitations. Hence, optimizations are hard to implement. The Spark team has commented on this, see <http://databricks.com/blog/2014/03/26/Spark-SQL-manipulating-structured-data-using-Spark.html>

Performance

Full dump to disk
between jobs.

Wednesday, April 23, 14

Sequencing jobs wouldn't be so bad if the "system" was smart enough to cache data in memory. Instead, each job dumps everything to disk, then the next job reads it back in again. This makes iterative algorithms particularly painful.



Enter
Spark
spark.apache.org

Cluster Computing

Can be run in:

- YARN (Hadoop 2)
- Mesos (Cluster management)
- EC2
- Standalone (“local” mode)



Compute Model

Fine-grained “combinators”
for composing algorithms.



Wednesday, April 23, 14

Once you learn the core set of primitives, it's easy to compose non-trivial algorithms with little code.

Compute Model

RDDs:
Resilient,
Distributed
Datasets



Wednesday, April 23, 14

RDDs shard the data over a cluster, like a virtualized, distributed collection (analogous to HDFS). They support intelligent caching, which means no naive flushes of massive datasets to disk. This feature alone allows Spark jobs to run 10-100x faster than comparable MapReduce jobs! The “resilient” part means they will reconstitute shards lost due to process/server crashes.

Compute Model

Written in Scala,
with Java and Python APIs.



Wednesday, April 23, 14

Once you learn the core set of primitives, it's easy to compose non-trivial algorithms with little code.



Inverted Index in MapReduce (Java).

Wednesday, April 23, 14

Let's see an actual implementation of the inverted index. First, a Hadoop MapReduce (Java) version, adapted from <https://developer.yahoo.com/hadoop/tutorial/module4.html#solution> It's about 90 lines of code, but I reformatted to fit better. This is also a slightly simpler version than the one I diagrammed. It doesn't record a count of each word in a document; it just writes (word,doc-title) pairs out of the mappers and the final (word,list) output by the reducers just has a list of documentations, hence repeats. A second job would be necessary to count the repeats.

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf =
            new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
    }
}
```

Wednesday, April 23, 14

I've shortened the original code a bit, e.g., using * import statements instead of separate imports for each class.

I'm not going to explain every line ... nor most lines.

Everything is in one outer class. We start with a main routine that sets up the job. Lotta boilerplate...

I used yellow for method calls, because methods do the real work!! But notice that the functions in this code don't really do a whole lot...

```
JobClient client = new JobClient();
JobConf conf =
    new JobConf(LineIndexer.class);

conf.setJobName("LineIndexer");
conf.setOutputKeyClass(Text.class);
conf.setOutputValueClass(Text.class);
FileInputFormat.addInputPath(conf,
    new Path("input"));
FileOutputFormat.setOutputPath(conf,
    new Path("output"));
conf.setMapperClass(
    LineIndexMapper.class);
conf.setReducerClass(
    LineIndexReducer.class);

client.setConf(conf);
```

```
    LineIndexMapper.class);  
    conf.setReducerClass(  
        LineIndexReducer.class);  
  
    client.setConf(conf);  
  
    try {  
        JobClient.runJob(conf);  
    } catch (Exception e) {  
        e.printStackTrace();  
    }  
}  
  
public static class LineIndexMapper  
    extends MapReduceBase  
    implements Mapper<LongWritable, Text,  
              Text, Text> {
```

```
public static class LineIndexMapper  
extends MapReduceBase  
implements Mapper<LongWritable, Text,  
Text, Text> {  
private final static Text word =  
new Text();  
private final static Text location =  
new Text();  
  
public void map(  
LongWritable key, Text val,  
OutputCollector<Text, Text> output,  
Reporter reporter) throws IOException {  
  
FileSplit fileSplit =  
(FileSplit)reporter.getInputSplit();  
String fileName -
```

```
FileSplit fileSplit =
  (FileSplit)reporter.getInputSplit();
String fileName =
  fileSplit.getPath().getName();
location.set(fileName);

String line = val.toString();
StringTokenizer itr = new
  StringTokenizer(line.toLowerCase());
while (itr.hasMoreTokens()) {
  word.set(itr.nextToken());
  output.collect(word, location);
}
}
```

```
public static class LineIndexReducer  
extends MapReduceBase  
implements Reducer<Text, Text,  
Text, Text> {  
    public void reduce(Text key,  
                      Iterator<Text> values,  
                      OutputCollector<Text, Text> output,  
                      Reporter reporter) throws IOException {  
        boolean first = true;  
        StringBuilder toReturn =  
            new StringBuilder();  
        while (values.hasNext()) {  
            if (!first)  
                toReturn.append(", ");  
            first=false;  
            toReturn.append(  
                values.next().toString());  
        }  
        output.collect(key, toReturn);  
    }  
}
```

Wednesday, April 23, 14

The reducer class, LineIndexReducer, with the reduce method that is called for each key and a list of values for that key. The reducer is stupid; it just reformats the values collection into a long string and writes the final (word,list-string) output.

```
boolean first = true;
StringBuilder toReturn =
    new StringBuilder();
while (values.hasNext()) {
    if (!first)
        toReturn.append(", ");
    first=false;
    toReturn.append(
        values.next().toString());
}
output.collect(key,
    new Text(toReturn.toString()));
}
}
```

Altogether

```
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LineIndexer {

    public static void main(String[] args) {
        JobClient client = new JobClient();
        JobConf conf =
            new JobConf(LineIndexer.class);

        conf.setJobName("LineIndexer");
        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(Text.class);
        FileInputFormat.addInputPath(conf,
            new Path("input"));
        FileOutputFormat.setOutputPath(conf,
            new Path("output"));
        conf.setMapperClass(
            LineIndexMapper.class);
        conf.setReducerClass(
            LineIndexReducer.class);

        client.setConf(conf);

        try {
            JobClient.runJob(conf);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    public static class LineIndexMapper
        extends MapReduceBase
        implements Mapper<LongWritable, Text,
                    Text, Text> {
        private final static Text word =
            new Text();
        private final static Text location =
            new Text();

        public void map(
            LongWritable key, Text val,
            OutputCollector<Text, Text> output,
            Reporter reporter) throws IOException {

            FileSplit fileSplit =
                (FileSplit)reporter.getInputSplit();
            String fileName =
                fileSplit.getPath().getName();
            location.set(fileName);

            String line = val.toString();
            StringTokenizer itr = new
                StringTokenizer(line.toLowerCase());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                output.collect(word, location);
            }
        }
    }

    public static class LineIndexReducer
        extends MapReduceBase
        implements Reducer<Text, Text,
                           Text, Text> {
        public void reduce(Text key,
                          Iterator<Text> values,
                          OutputCollector<Text, Text> output,
                          Reporter reporter) throws IOException {
            boolean first = true;
            StringBuilder toReturn =
                new StringBuilder();
            while (values.hasNext()) {
                if (!first)
                    toReturn.append(", ");
                first=false;
                toReturn.append(
                    values.next().toString());
            }
            output.collect(key,
                new Text(toReturn.toString()));
        }
    }
}
```

Wednesday, April 23, 14

The whole shebang (6pt. font)



Inverted Index in Spark (Scala).

Wednesday, April 23, 14

This code is approximately 45 lines, but it does more than the previous Java example, it implements the original inverted index algorithm I diagrammed where word counts are computed and included in the data.

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
          text.split("\\W+").map(s =>
```

Wednesday, April 23, 14

It starts with imports, then declares a singleton object (a first-class concept in Scala), with a main routine (as in Java).

The methods are colored yellow again. Note this time how dense with meaning they are this time.

```
val sc = new SparkContext(  
  "local", "Inverted Index")  
  
sc.textFile("data/crawl")  
.map { line =>  
  val array = line.split("\t", 2)  
  (array(0), array(1))  
}  
.flatMap {  
  case (path, text) =>  
    text.split("""\w+""") map {  
      word => (word, path)  
    }  
}  
.map {  
  case (w, p) => ((w, p), 1)  
}
```

Wednesday, April 23, 14

You begin the workflow by declaring a SparkContext (in “local” mode, in this case). The rest of the program is a sequence of function calls, analogous to “pipes” we connect together to perform the data flow.

Next we read one or more text files. If “data/crawl” has 1 or more Hadoop-style “part-NNNNN” files, Spark will process all of them (in parallel if running a distributed configuration; they will be processed synchronously in local mode).

```
.map { line =>
    val array = line.split("\t", 2)
    (array(0), array(1))
}
.flatMap {
    case (path, text) =>
    text.split("""\w+""") map {
        word => (word, path)
    }
}
.map {
    case (w, p) => ((w, p), 1)
}
.reduceByKey {
    case (n1, n2) => n1 + n2
}
.map {
```

Wednesday, April 23, 14

sc.textFile returns an RDD with a string for each line of input text. So, the first thing we do is map over these strings to extract the original document id (i.e., file name), followed by the text in the document, all on one line. We assume tab is the separator. "(array(0), array(1))" returns a two-element "tuple". Think of the output RDD has having a schema "String fileName, String text".

```
}

.flatMap {
    case (path, text) =>
        text.split("""\w+""") map {
            word => (word, path)
        }
}
.map {
    case (w, p) => ((w, p), 1)
}
.reduceByKey {
    case (n1, n2) => n1 + n2
}
.map {
    case ((w, p), n) => (w, (p, n))
}
.groupBy {
```

*Beautiful,
no?*

Wednesday, April 23, 14

flatMap maps over each of these 2-element tuples. We split the text into words on non-alphanumeric characters, then output collections of word (our ultimate, final “key”) and the path. Each line is converted to a collection of (word,path) pairs, so flatMap converts the collection of collections into one long “flat” collection of (word,path) pairs.

Then we map over these pairs and add a single count of 1.

$$\nabla \cdot D = \rho$$

$$\nabla \cdot B = 0$$

$$\nabla \times E = -\frac{\partial B}{\partial t}$$

$$\nabla \times H = J + \frac{\partial D}{\partial t}$$

Wednesday, April 23, 14

Another example of a beautiful and profound DSL, in this case from the world of Physics: Maxwell's equations: <http://upload.wikimedia.org/wikipedia/commons/c/c4/Maxwell%27sEquations.svg>

```
}

.reduceByKey {
    case (n1, n2) => n1 + n2
}

.map {
    case ((w, p), n) => (w, (p, n))
}

(groupBy {
    case (w, (p, n)) => w ...
}

.map {
    case (w, seq) =>
        val seq2 = seq map {
            case (_, (p, n)) => (p, n)
        }
        (w, seq2.mkString("", ""))
}
```

Wednesday, April 23, 14

reduceByKey does an implicit “group by” to bring together all occurrences of the same (word, path) and then sums up their counts.
Note the input to the next map is now ((word, path), n), where n is now ≥ 1 . We transform these tuples into the form we actually want, (word, (path, n)).

```

}
    .groupBy {
        case (w, (p, n)) => w
    }      (word, seq((word, (path1, n1)), (word, (path2, n2)), ...))
    .map {
        case (w, seq) =>
        val seq2 = seq map {
            case (_, (p, n)) => (p, n)
        }
        (w, seq2.mkString(", "))
    }          (word, "(path1, n1), (path2, n2), ...")
    .saveAsTextFile(argz.outpath)

    sc.stop()
}
}

```

Wednesday, April 23, 14

Now we do an explicit group by to bring all the same words together. The output will be (word, (word, (path1, n1)), (word, (path2, n2)), ...). The last map removes the redundant “word” values in the sequences of the previous output. It outputs the sequence as a final string of comma-separated (path,n) pairs.

We finish by saving the output as text file(s) and stopping the workflow.

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object InvertedIndex {
  def main(args: Array[String]) = {

    val sc = new SparkContext(
      "local", "Inverted Index")

    sc.textFile("data/crawl")
      .map { line =>
        val array = line.split("\t", 2)
        (array(0), array(1))
      }
      .flatMap {
        case (path, text) =>
        text.split("""\W+""") map {
          word => (word, path)
        }
      }
      .map {
        case (w, p) => ((w, p), 1)
      }
      .reduceByKey {
        case (n1, n2) => n1 + n2
      }
      .map {
        case ((w, p), n) => (w, (p, n))
      }
      .groupByKey {
        case (w, (p, n)) => w
      }
      .map {
        case (w, seq) =>
        val seq2 = seq map {
          case (_, (p, n)) => (p, n)
        }
        (w, seq2.mkString(", "))
      }
      .saveAsTextFile(argz.outpath)

    sc.stop()
  }
}
```

Altogether

```
sc.textFile("data/crawl")
.map { line =>
  val array = line.split("\t", 2)
  (array(0), array(1))
}
.flatMap {
  case (path, text) =>
  text.split("""\W+""") map {
    word => (word, path)
  }
}
.map {
  case (w, p) => ((w, p), 1)
}
.reduceByKey{
  case (n1, n2) => n1 + n2
}
```

*Powerful
Combinators!*

Wednesday, April 23, 14

I've shortened the original code a bit, e.g., using * import statements instead of separate imports for each class.
I'm not going to explain every line ... nor most lines.
Everything is in one outer class. We start with a main routine that sets up the job. Lotta boilerplate...



The Spark version
took me ~30 minutes
to write.

Wednesday, April 23, 14

Once you learn the core primitives I used, and a few tricks for manipulating the RDD tuples, you can very quickly build complex algorithms for data processing!
The Spark API allowed us to focus almost exclusively on the “domain” of data transformations, while the Java MapReduce version (which does less), forced tedious attention to infrastructure mechanics.

SQL!

Shark:
Hive (SQL query tool)
ported to Spark.





*Use a SQL query when
you can!!*

```
CREATE EXTERNAL TABLE stocks (
    symbol STRING,
    ymd STRING,
    price_open STRING,
    price_close STRING,
    shares_traded INT)
LOCATION "hdfs://data/stocks";
```

```
-- Year-over-year average closing price.
SELECT year(ymd), avg(price_close)
FROM stocks
WHERE symbol = 'AAPL'
GROUP BY year(ymd);
```

Hive Query Language.

Shark

~10-100x the performance of
Hive, due to in-memory
caching of RDDs & better
Spark abstractions.



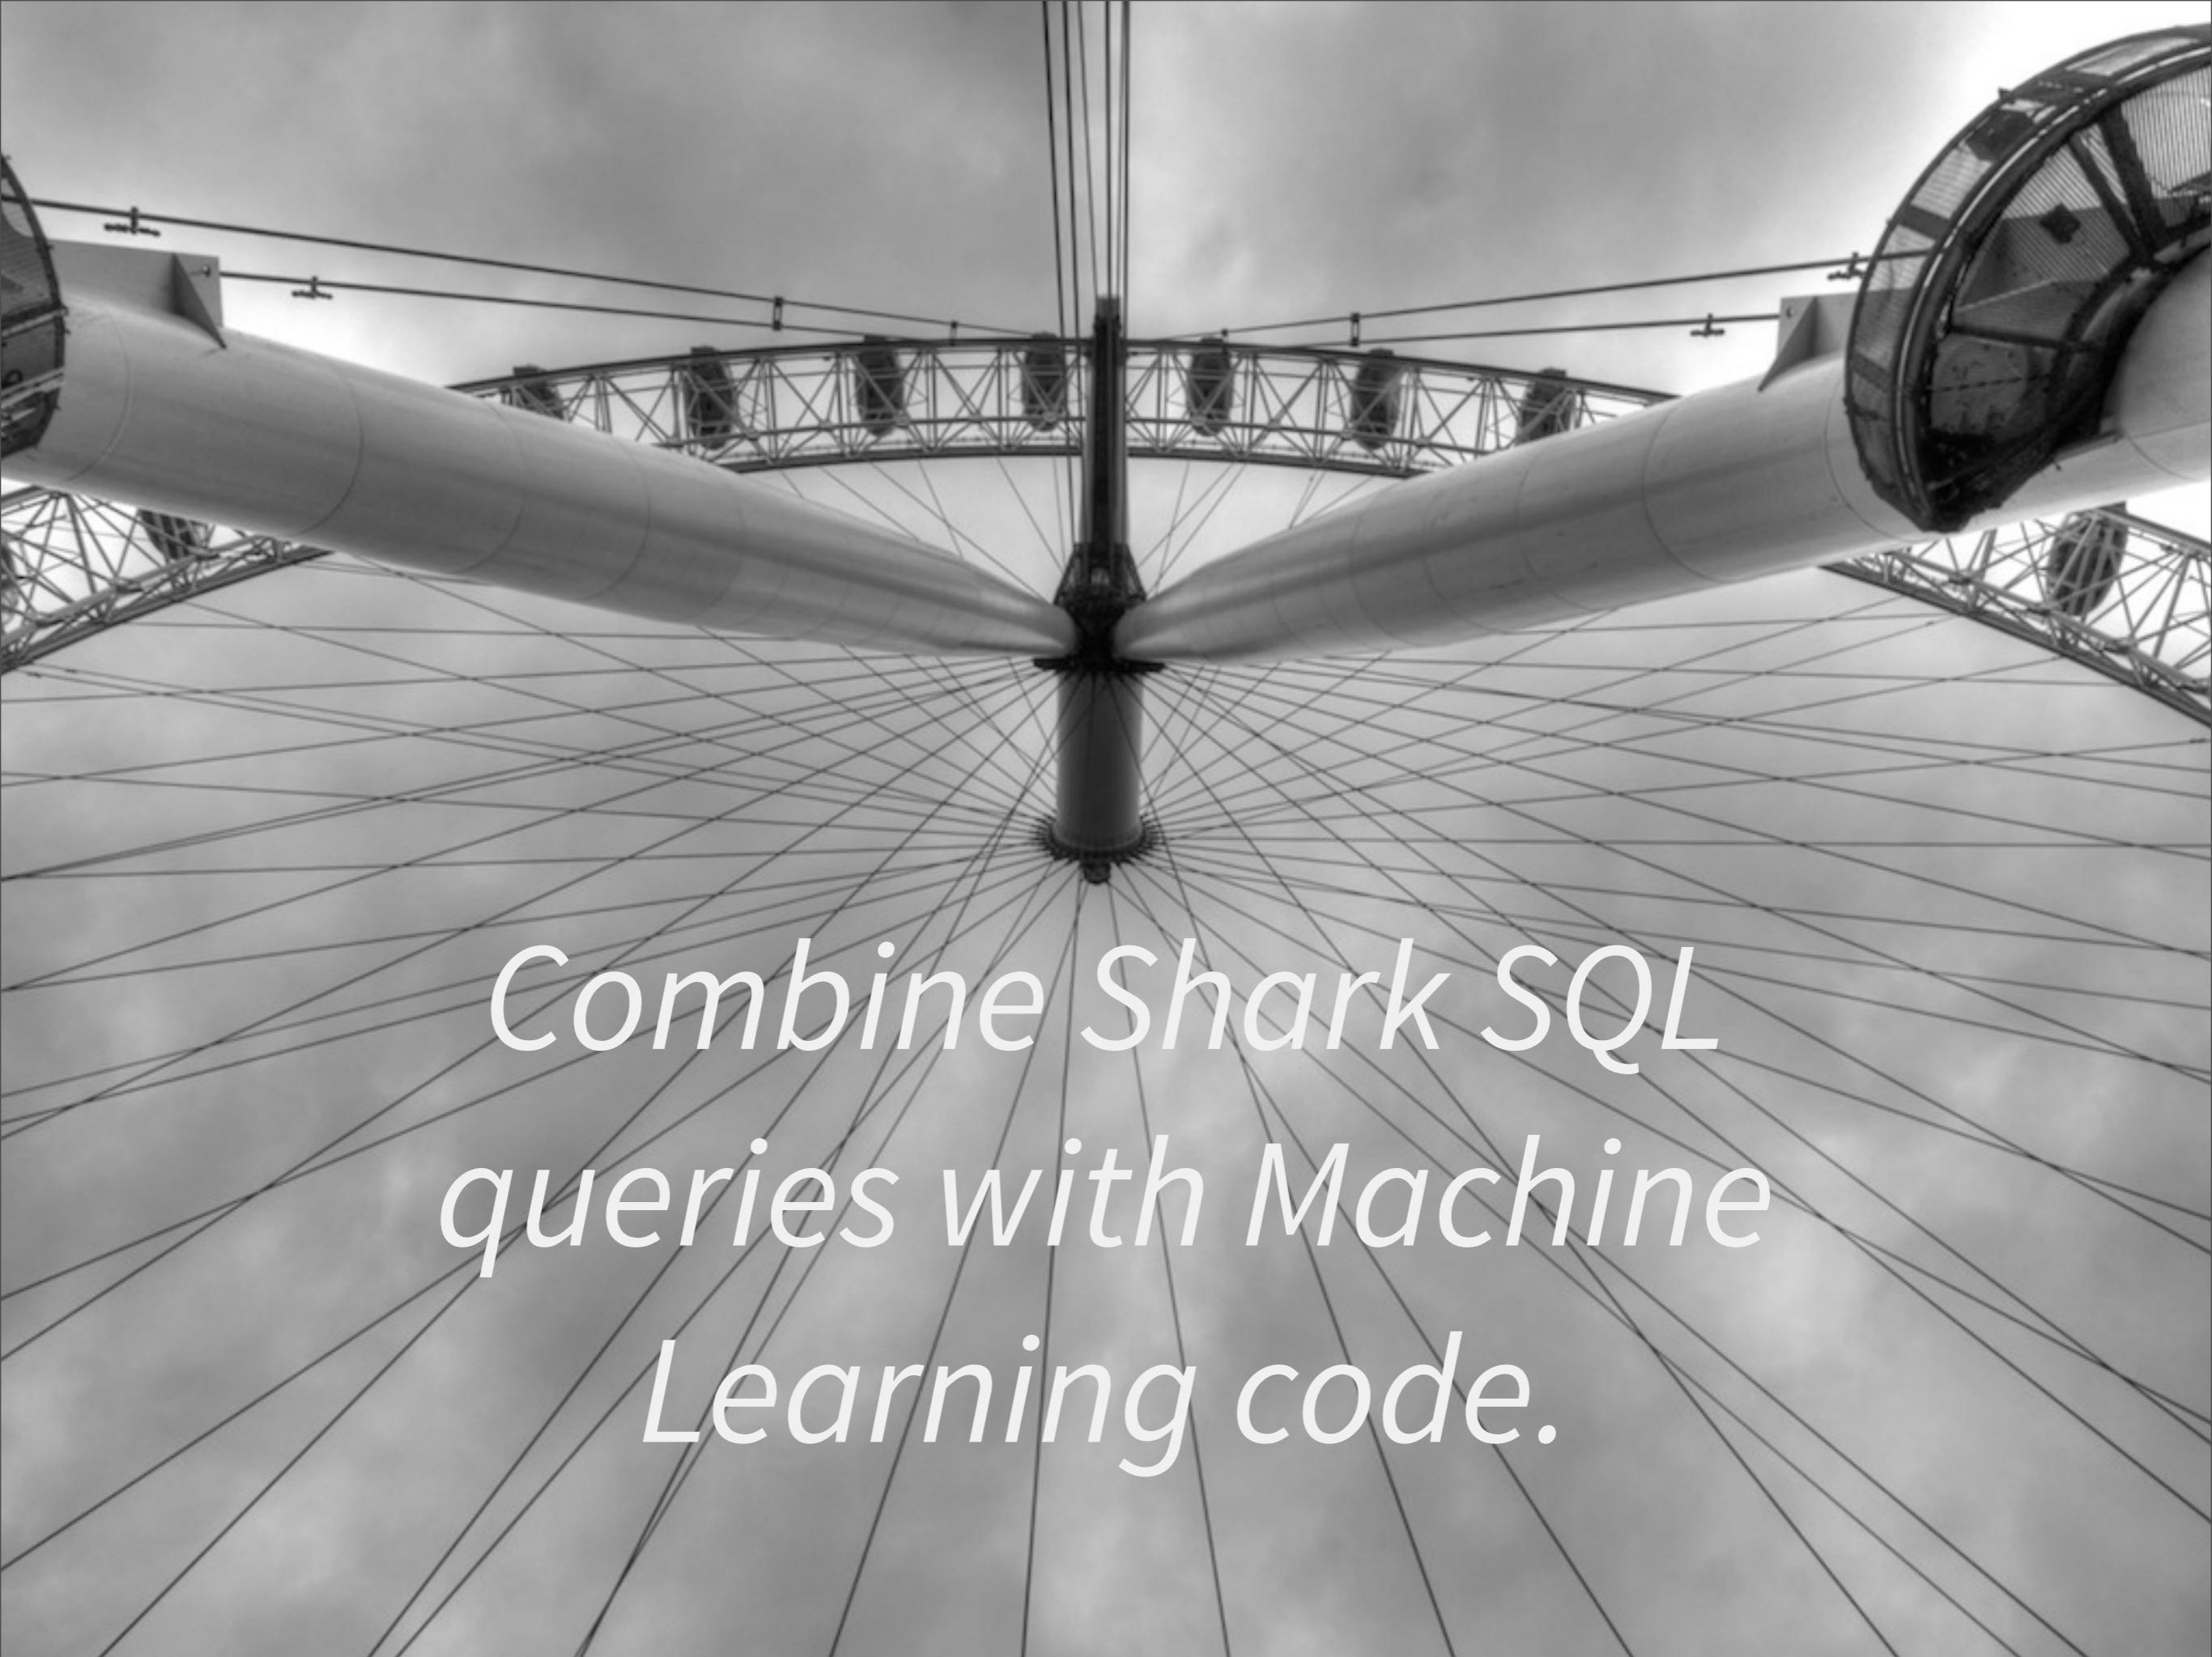
Did I mention SQL?

Spark SQL:
Next generation
SQL query tool/API.



Wednesday, April 23, 14

A new query optimizer. It will become the basis for Shark internals, replacing messy Hive code that is hard to reason about, extend, and debug.



Combine Shark SQL queries with Machine Learning code.

Wednesday, April 23, 14

We'll use the Spark "MLlib" in the example, then return to it in a moment.

```
CREATE TABLE Users(  
  userId INT,  
  name STRING,  
  email STRING,  
  age INT,  
  latitude DOUBLE,  
  longitude DOUBLE,  
  subscribed BOOLEAN);
```

Hive/Shark table definitions (not Scala).

```
CREATE TABLE Events(  
  userId INT,  
  action INT);
```

Wednesday, April 23, 14

This example adapted from the following blog post announcing Spark SQL:

<http://databricks.com/blog/2014/03/26/Spark-SQL-manipulating-structured-data-using-Spark.html>

Assume we have these Hive/Shark tables, with data about our users and events that have occurred.

```
val trainingDataTable = sql("""  
SELECT e.action, u.age,  
u.latitude, u.longitude  
FROM Users u  
JOIN Events e  
ON u.userId = e.userId""")
```

Spark Scala.

```
val trainingData =  
trainingDataTable map { row =>  
val features =  
  Array[Double](row(1), row(2), row(3))  
LabeledPoint(row(0), features)  
}  
  
val model =  
new LogisticRegressionWithSGD()
```

Wednesday, April 23, 14

Here is some Spark (Scala) code with an embedded SQL/Shark query that joins the Users and Events tables. The "'''...''' string allows embedded line feeds.

The “sql” function returns an RDD, which we then map over to create LabeledPoints, an object used in Spark’s MLlib (machine learning library) for a recommendation engine. The “label” is the kind of event and the user’s age and lat/long coordinates are the “features” used for making recommendations. (E.g., if you’re 25 and near a certain location in the city, you might be interested a nightclub near by...)

```
val model =  
  new LogisticRegressionWithSGD()  
  .run(trainingData)  
  
val allCandidates = sql("""  
  SELECT userId, age, latitude, longitude  
  FROM Users  
  WHERE subscribed = FALSE""")  
  
case class Score(  
  userId: Int, score: Double)  
val scores = allCandidates map { row =>  
  val features =  
    Array[Double](row(1), row(2), row(3))  
  Score(row(0), model.predict(features))  
}  
}
```

Wednesday, April 23, 14

Next we train the recommendation engine, using a “logistic regression” fit to the training data, where “stochastic gradient descent” (SGD) is used to train it. (This is a standard tool set for recommendation engines; see for example: <http://www.cs.cmu.edu/~wcohen/10-605/assignments/sgd.pdf>)

```
val allCandidates = sql("""  
    SELECT userId, age, latitude, longitude  
    FROM Users  
    WHERE subscribed = FALSE""")  
  
case class Score(  
    userId: Int, score: Double)  
val scores = allCandidates map { row =>  
    val features =  
        Array[Double](row(1), row(2), row(3))  
    Score(row(0), model.predict(features))  
}  
  
// Hive table  
scores.registerAsTable("Scores")
```

```
case class Score(  
    userId: Int, score: Double)  
val scores = allCandidates map { row =>  
    val features =  
        Array[Double](row(1), row(2), row(3))  
    Score(row(0), model.predict(features))  
}  
  
// Hive table  
scores.registerAsTable("Scores")  
  
val topCandidates = sql("""  
    SELECT u.name, u.email  
    FROM Scores s  
    JOIN Users u ON s.userId = u.userId  
    ORDER BY score DESC
```

Wednesday, April 23, 14

Declare a class to hold each user's "score" as produced by the recommendation engine and map the "all" query results to Scores.
Then "register" the scores RDD as a "Scores" table in Hive's metadata repository. This is equivalent to running a "CREATE TABLE Scores ..." command at the Hive/Shark prompt!

```
// Hive table  
scores.registerAsTable("Scores")  
  
val topCandidates = sql("""  
SELECT u.name, u.email  
FROM Scores s  
JOIN Users u ON s.userId = u.userId  
ORDER BY score DESC  
LIMIT 100""")
```

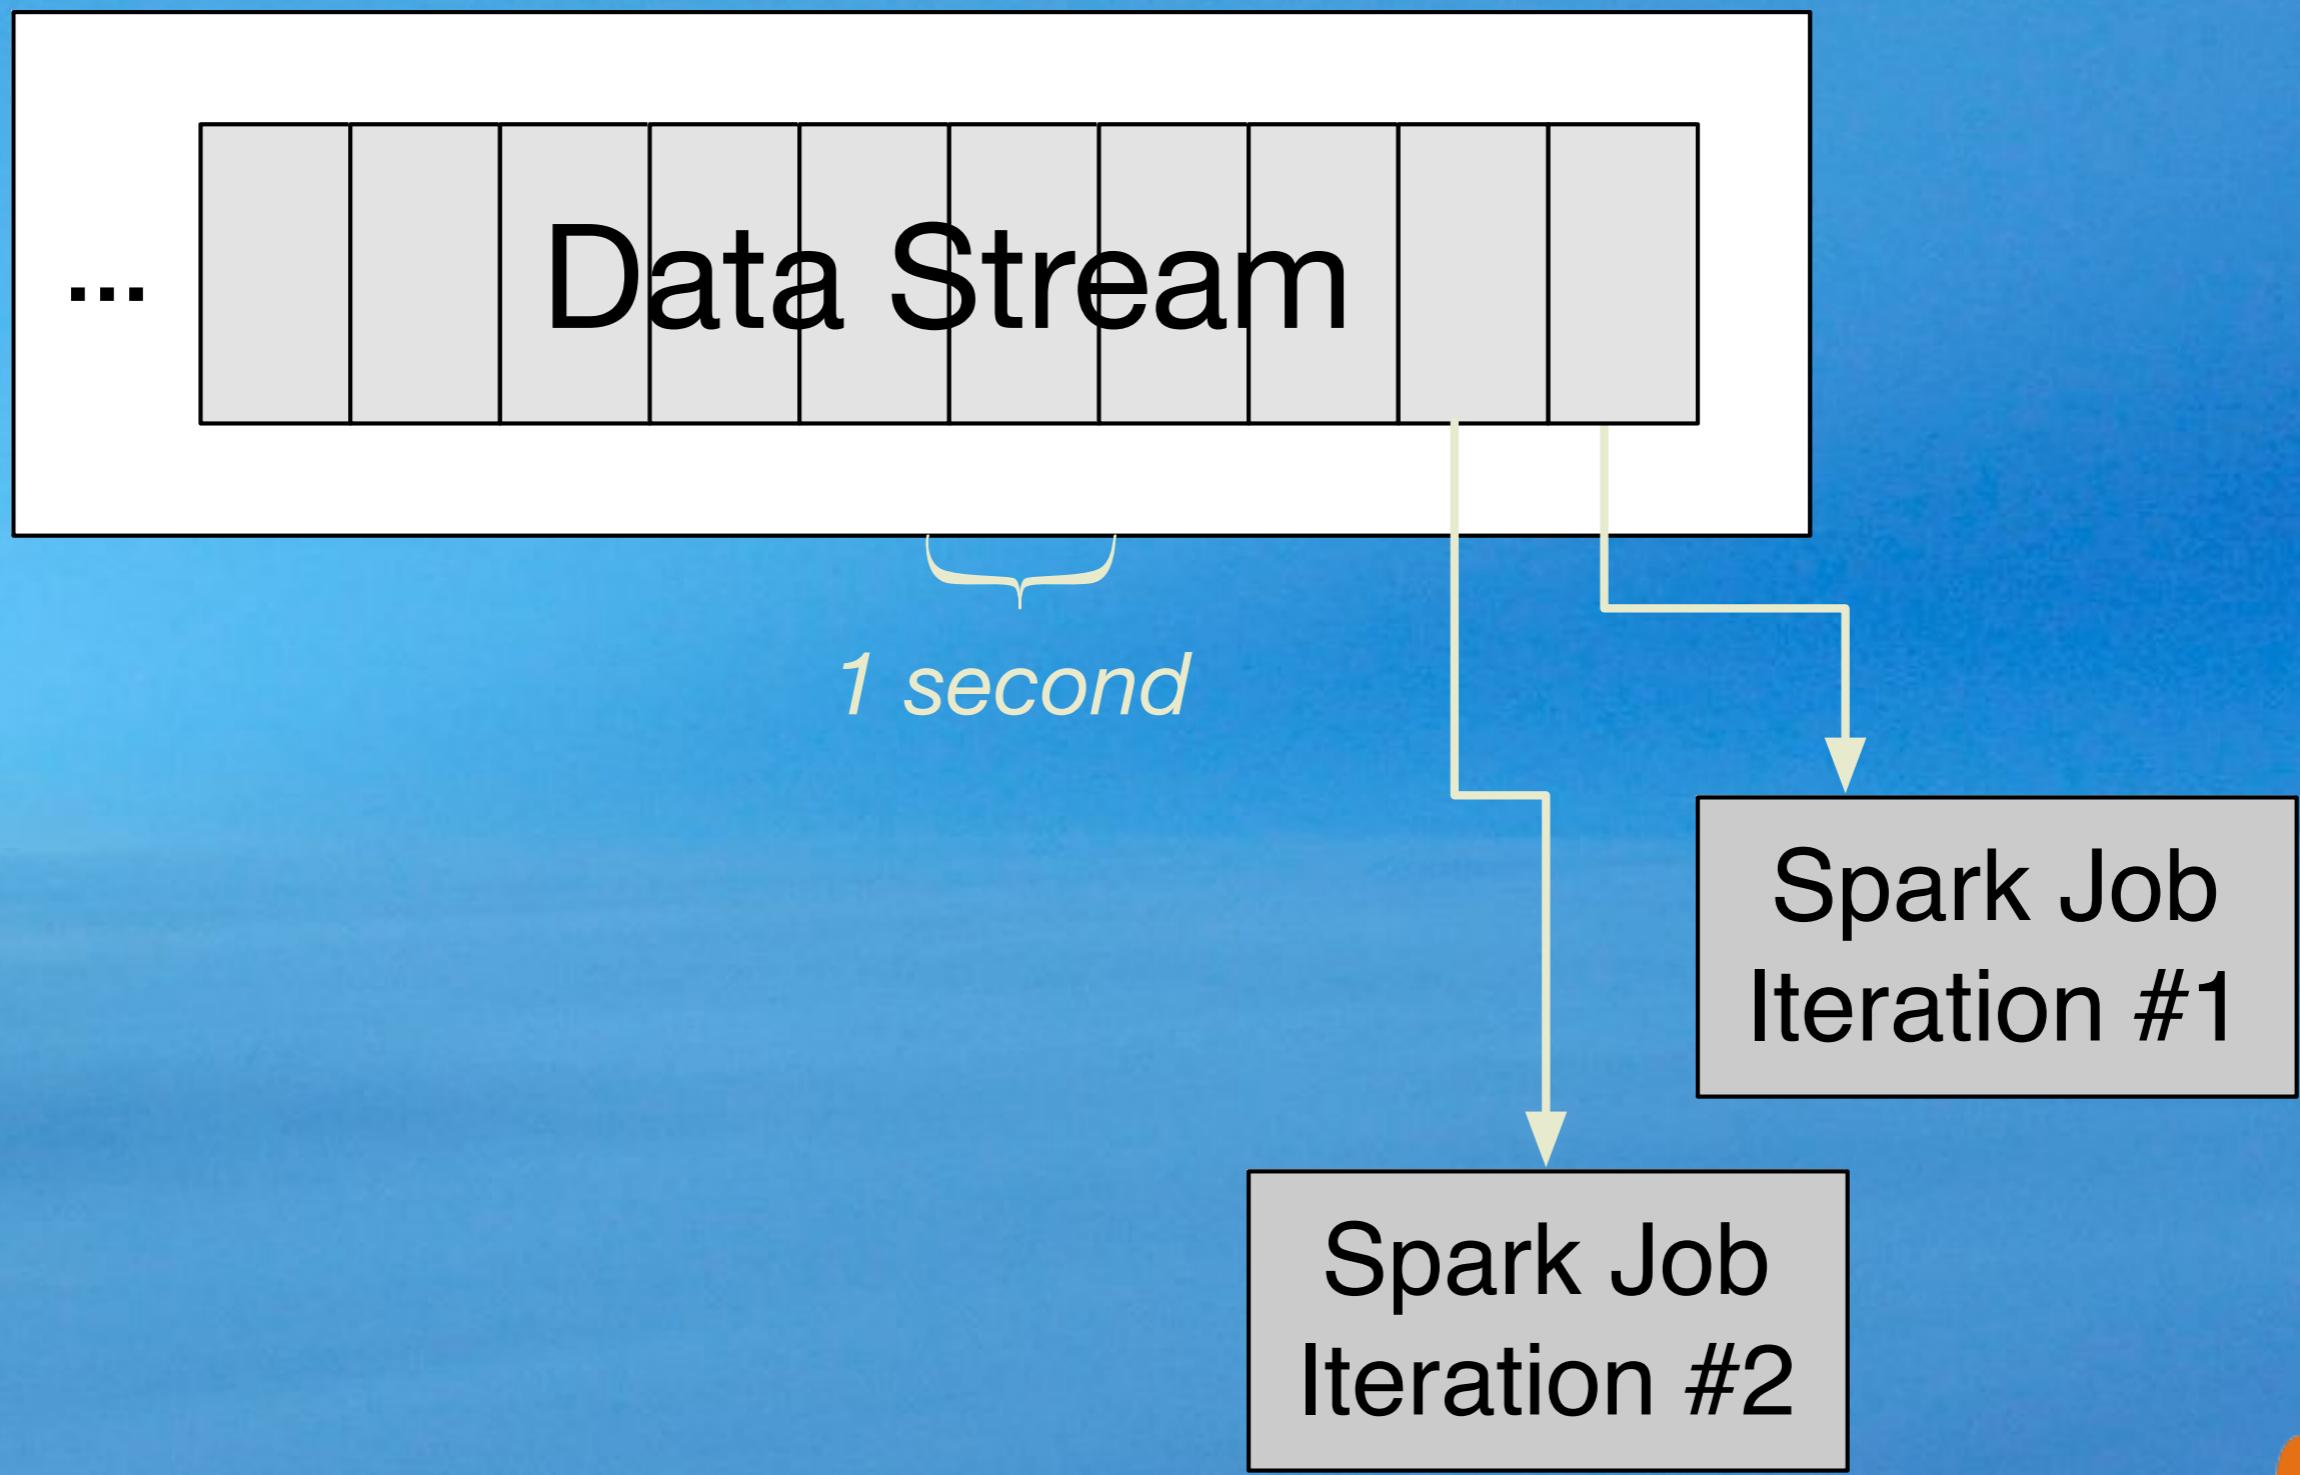
Cluster Computing

Spark Streaming:
Use the same abstractions for
real-time, event streaming.



Wednesday, April 23, 14

Once you learn the core set of primitives, it's easy to compose non-trivial algorithms with little code.



Spark 

Wednesday, April 23, 14

You can specify the granularity, such as all events in 1 second windows, then your Spark job is patched each window of data for processing.



Very similar code...

```
val sc = new SparkContext(...)  
val ssc = new StreamingContext(  
    sc, Seconds(1))  
  
// A DStream that will listen to server:port  
val lines =  
    ssc.socketTextStream(server, port)  
  
// Word Count...  
val words = lines flatMap {  
    line => line.split("""\W+""")  
}  
  
val pairs = words map (word => (word, 1))  
val wordCounts =  
    pairs reduceByKey ((n1, n2) => n1 + n2)
```

Wednesday, April 23, 14

This example adapted from the following page on the Spark website:

<http://spark.apache.org/docs/0.9.0/streaming-programming-guide.html#a-quick-example>

We create a StreamingContext that wraps a SparkContext (there are alternative ways to construct it...). It will “clump” the events into 1-second intervals. Next we setup a socket to stream text to us from another server and port (one of several ways to ingest data).

```
ssc.socketTextStream(server, port)
```

```
// Word Count...
val words = lines flatMap {
  line => line.split("""\W+""")
}

val pairs = words map (word => (word, 1))
val wordCounts =
  pairs reduceByKey ((n1, n2) => n1 + n2)

wordCount.print() // print a few counts...

ssc.start()
ssc.awaitTermination()
```

Wednesday, April 23, 14

Now the “word count” happens over each interval (aggregation across intervals is also possible), but otherwise it works like before.

Once we setup the flow, we start it and wait for it to terminate through some means, such as the server socket closing.

Cluster Computing

MLlib:
Machine learning library.



Wednesday, April 23, 14

Spark implements many machine learning algorithms, although a lot more or needed, compared to more mature tools like Mahout and libraries in Python and R.

MLlib

- Linear regression
- Binary classification
- Collaborative filtering
- Clustering
- Others...



Wednesday, April 23, 14

Not as full-featured as more mature toolkits, but the Mahout project has announced they are going to port their algorithms to Spark, which include powerful Mathematics, e.g., Matrix support libraries.

Cluster Computing

GraphX:
Graphical models
and algorithms.



Wednesday, April 23, 14

Some problems are more naturally represented as graphs.
Extends RDDs to support property graphs with directed edges.



Dean Wampler

@deanwampler

Functional Programming: I came for the concurrency, but I stayed for the data science.

Reply Delete ★ Favorite ... More

RETWEETS

6

FAVORITES

5



Wednesday, April 23, 14

Why is Spark so good (and Java MapReduce so bad)? Because fundamentally, data analytics is Mathematics and programming tools inspired by Mathematics - like Functional Programming - are ideal tools for working with data. This is why Spark code is so concise, yet powerful. This is why it is a great platform for performance optimizations. This is why Spark is a great platform for higher-level tools, like SQL, graphs, etc. Interest in FP started growing ~10 years ago as a tool to attack concurrency. I believe that data is now driving FP adoption even faster. I know many Java shops that switched to Scala when they adopted tools like Spark and Scalding (<https://github.com/twitter/scalding>).

Spark

A flexible, scalable distributed
compute platform with
concise, powerful APIs and
higher-order tools.

spark.apache.org



A photograph of the London skyline from across the River Thames. On the left, the London Eye Ferris wheel stands prominently. In the center, the green Westminster Bridge spans the river. To the right, the historic Palace of Westminster and its iconic Elizabeth Tower (Big Ben) are visible against a dramatic, cloudy sky.

Why Spark Is the Next Top (Compute) Model

@deanwampler
polyglotprogramming.com/talks