



Lightbend



Stream All the Things!

Architectures for Data that Never Ends

Dean Wampler, Ph.D.
dean@lightbend.com
[@deanwampler](https://twitter.com/deanwampler)

Free as in 

lightbend.com/fast-data-platform
(2nd Edition coming soon!)

O'REILLY®

Fast Data Architectures for Streaming Applications

Getting Answers Now from
Data Sets that Never End

Dean Wampler

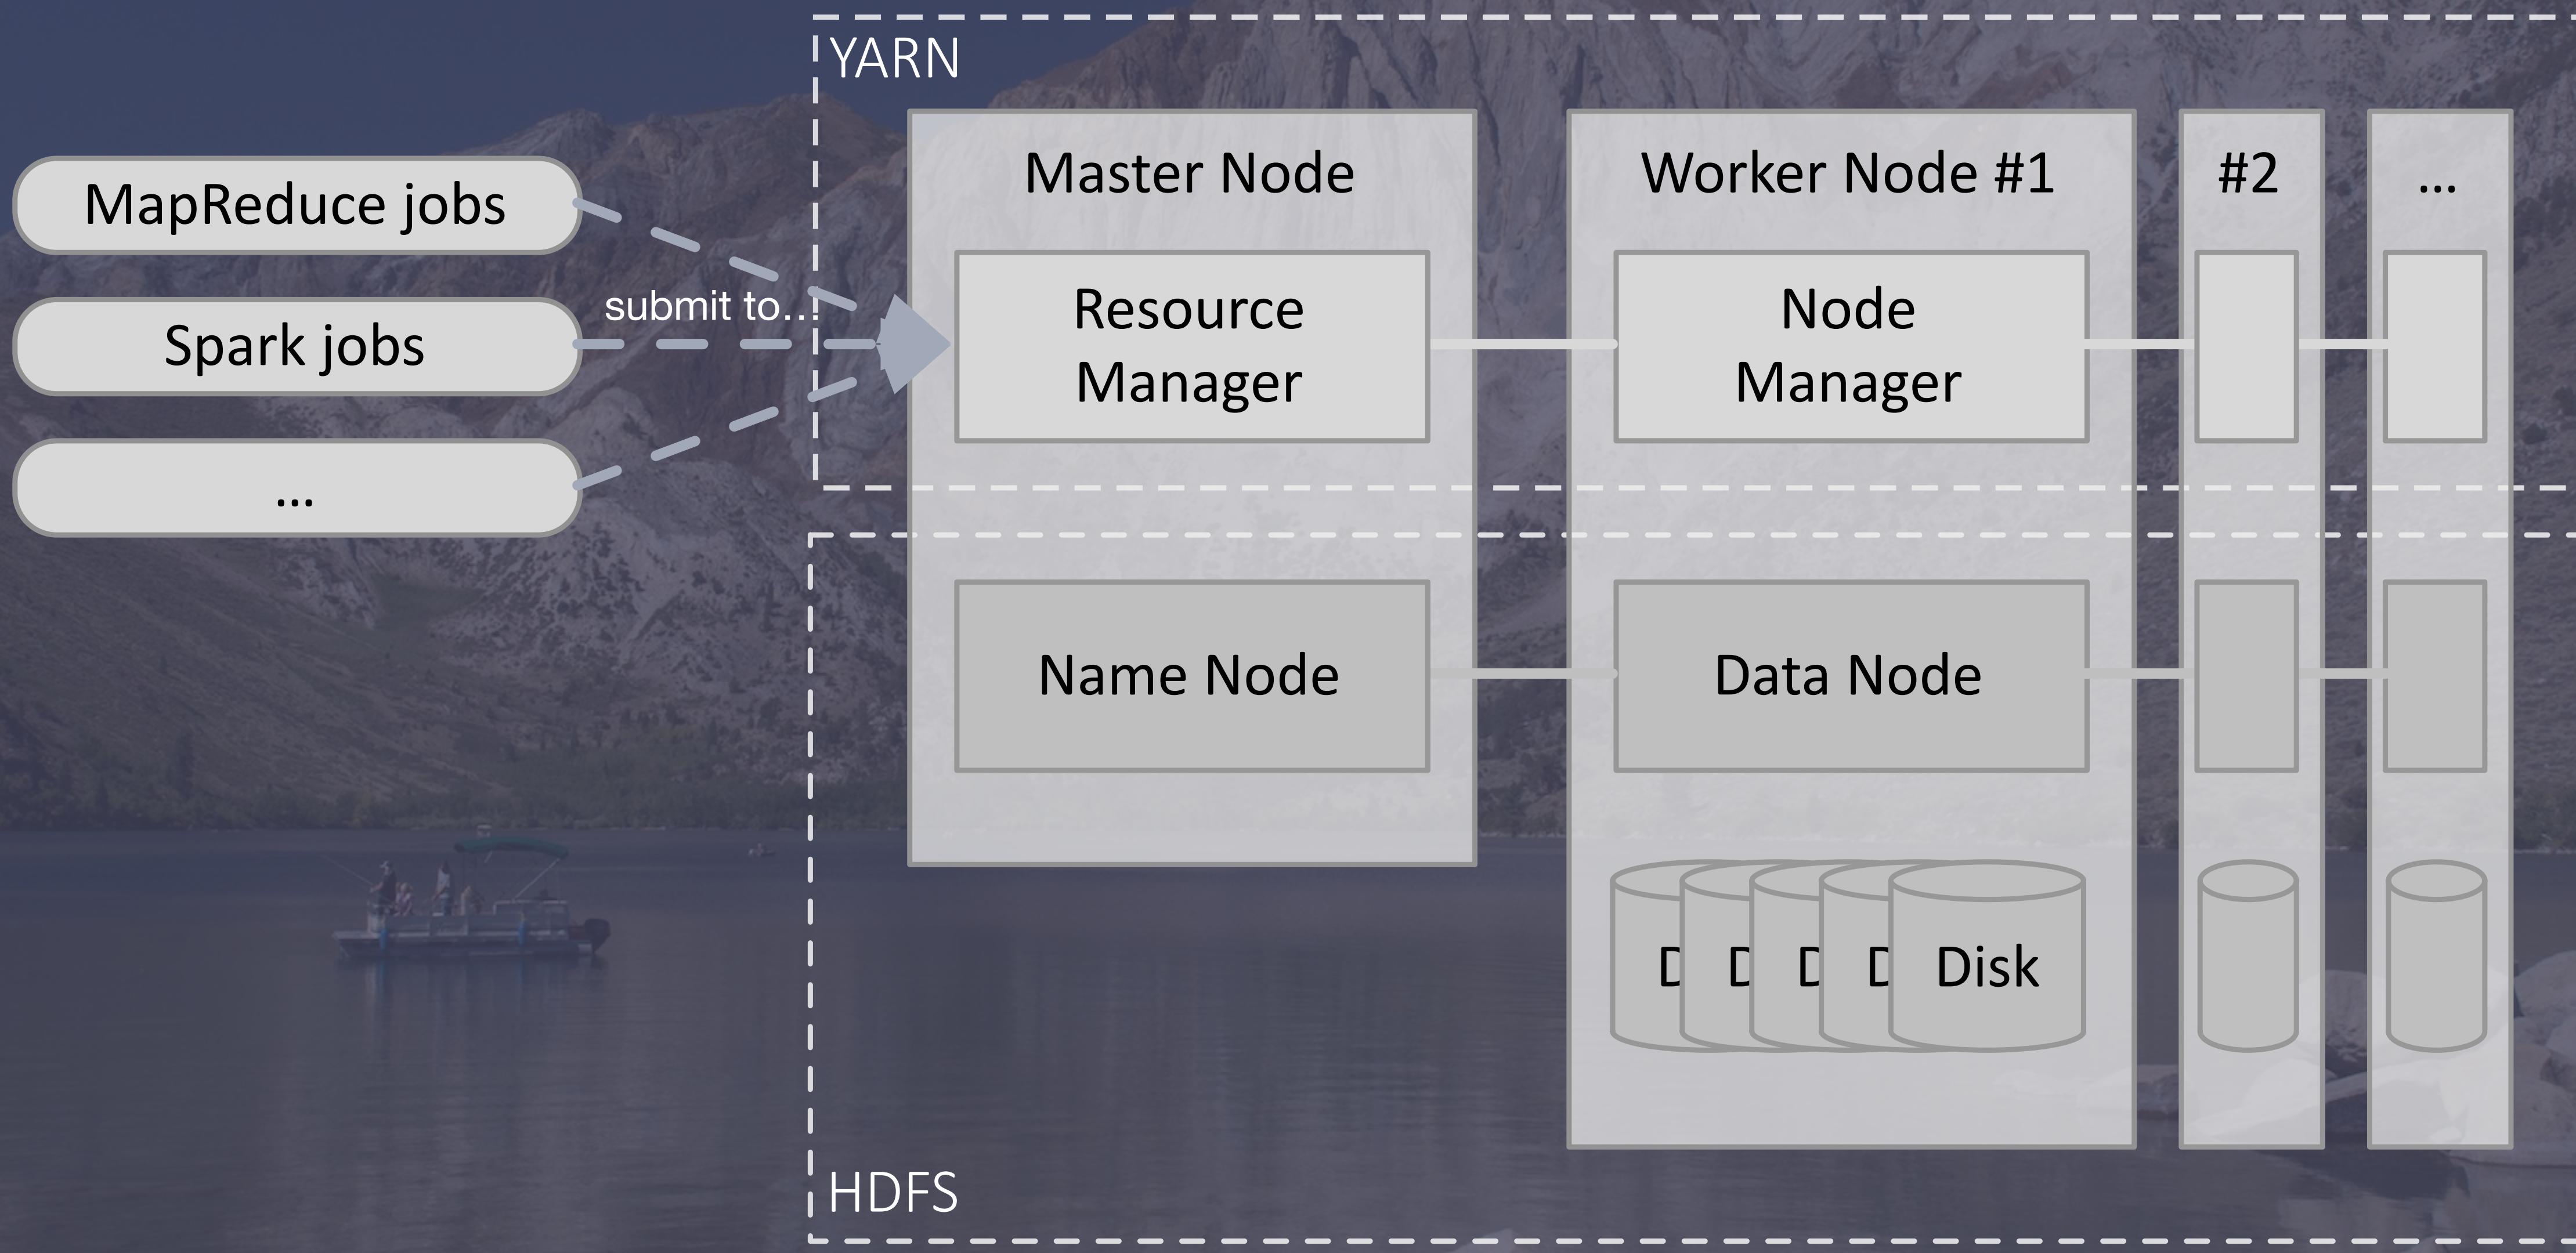
Compliments of
 Lightbend

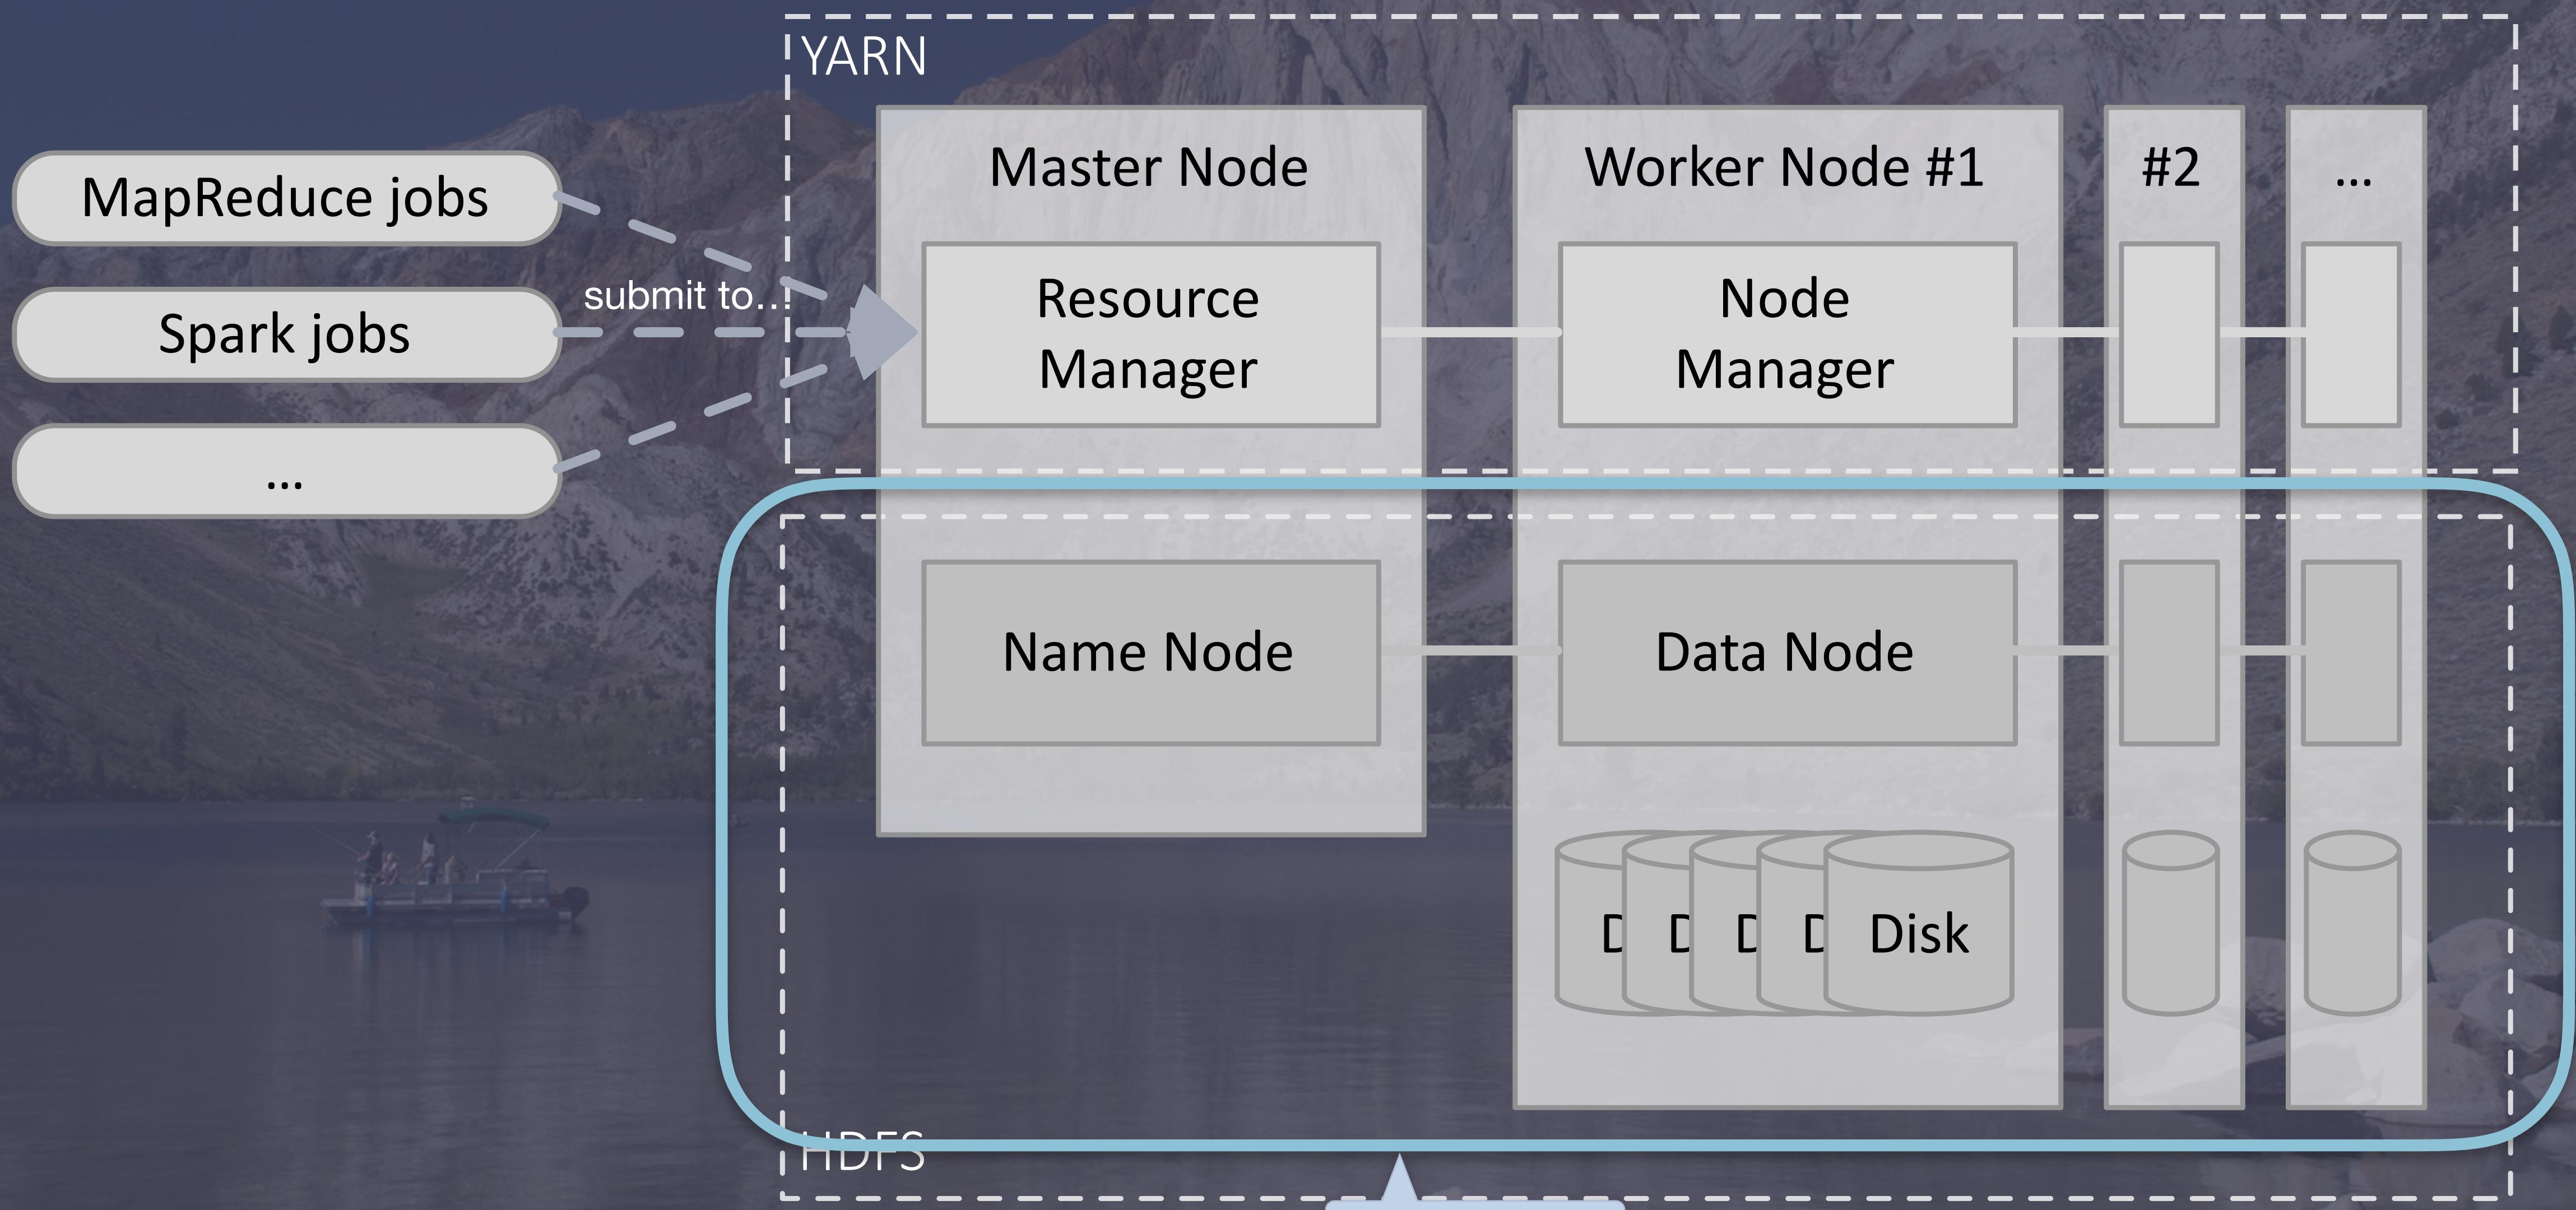
An aerial photograph showing a vast area of agricultural land. The fields are organized into a grid-like pattern, with various shades of green and brown indicating different crops or stages of cultivation. A network of blue waterways, including rivers and canals, cuts through the fields. The perspective is from above, looking down at the landscape.

Streaming
in Context...

A scenic landscape featuring a large, rugged mountain range with sharp peaks and vertical rock faces. In the foreground, there's a calm lake with a small, white pontoon boat containing two people. The shoreline in the bottom right corner is made of large, light-colored rocks. The sky is clear and blue.

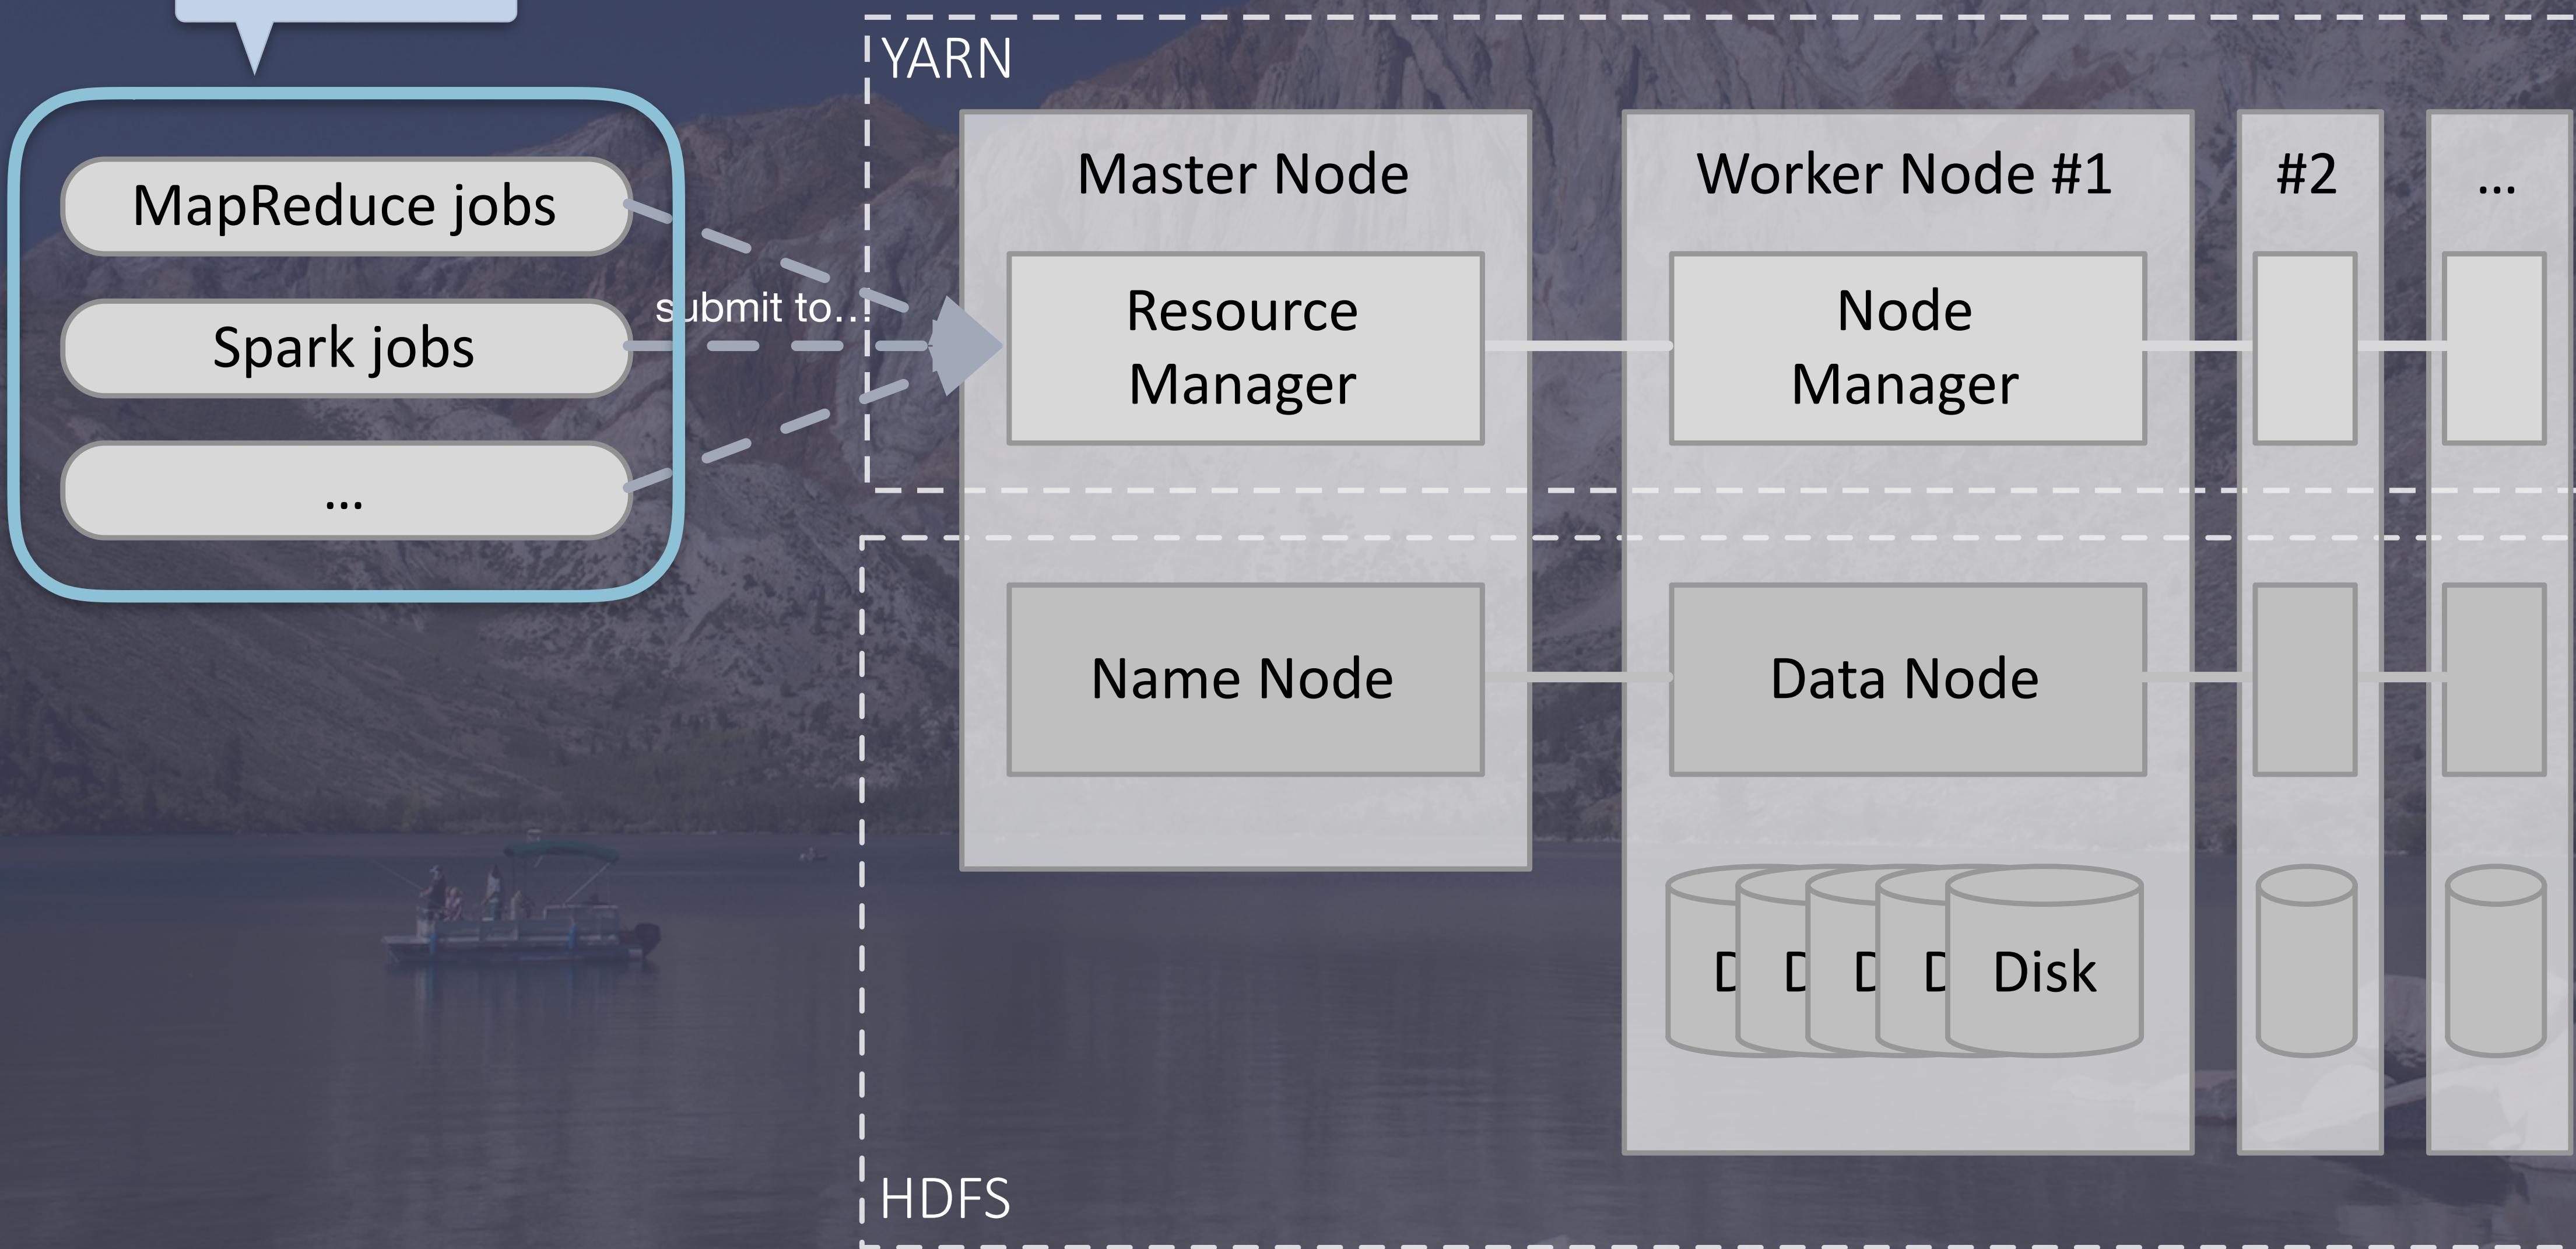
Hadoop: Classic Batch Architecture

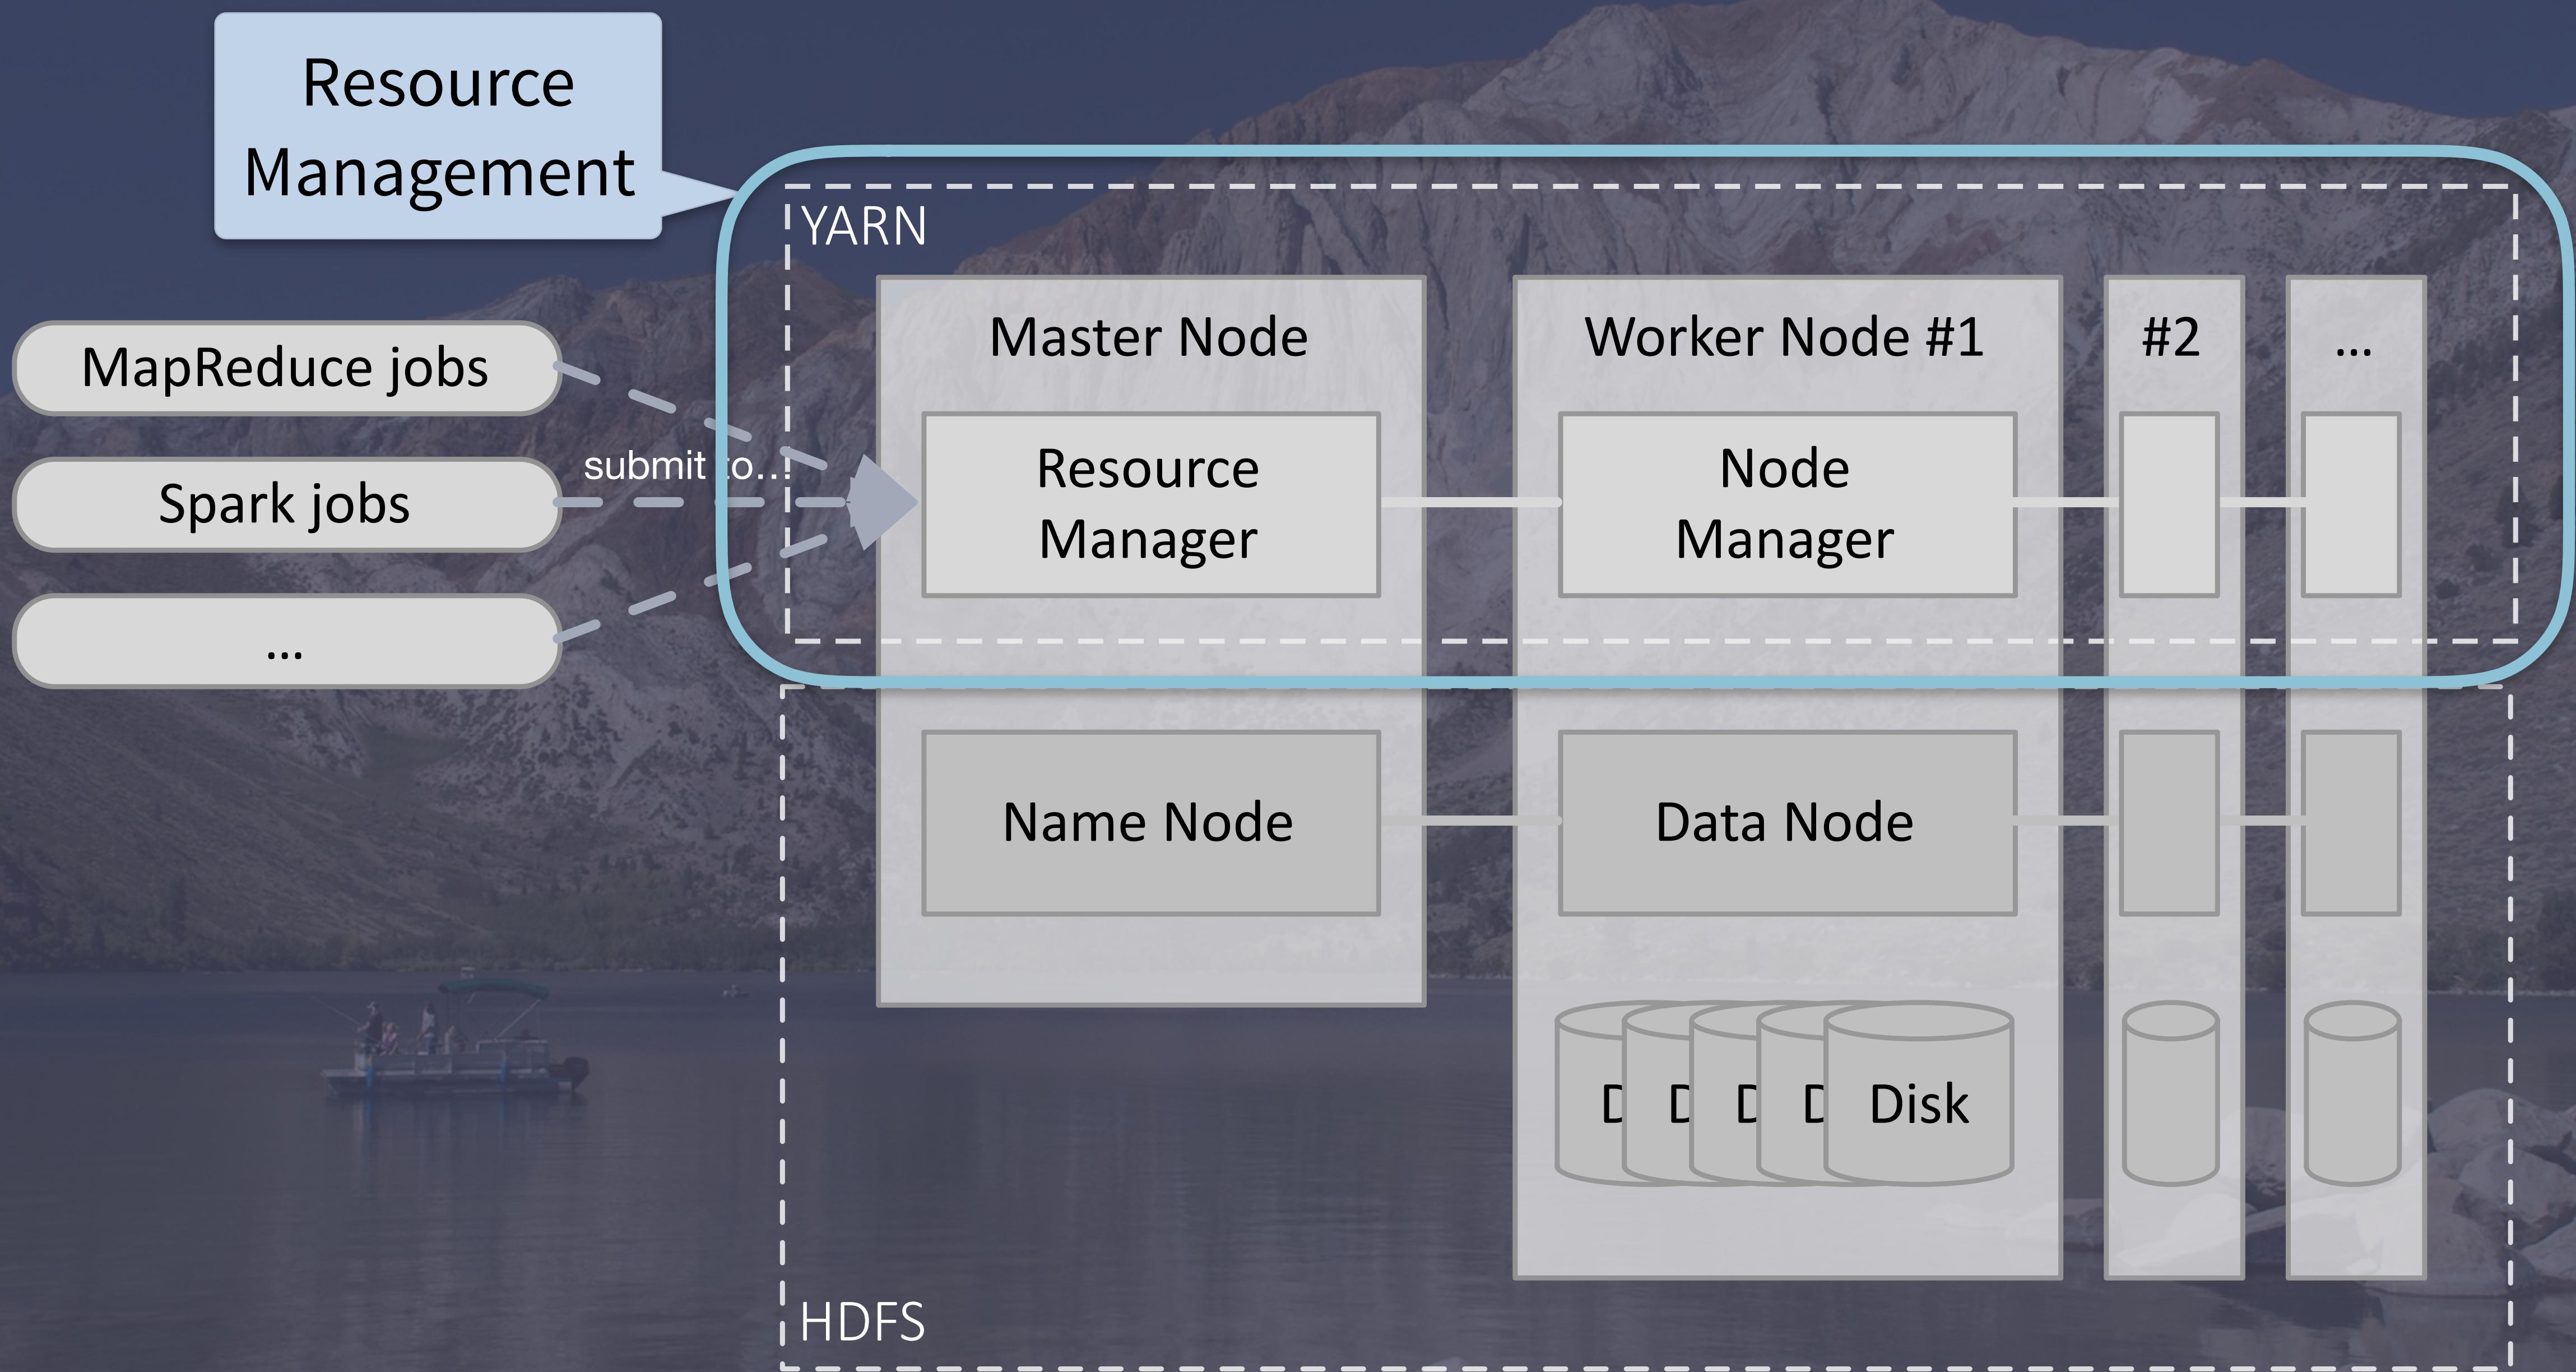




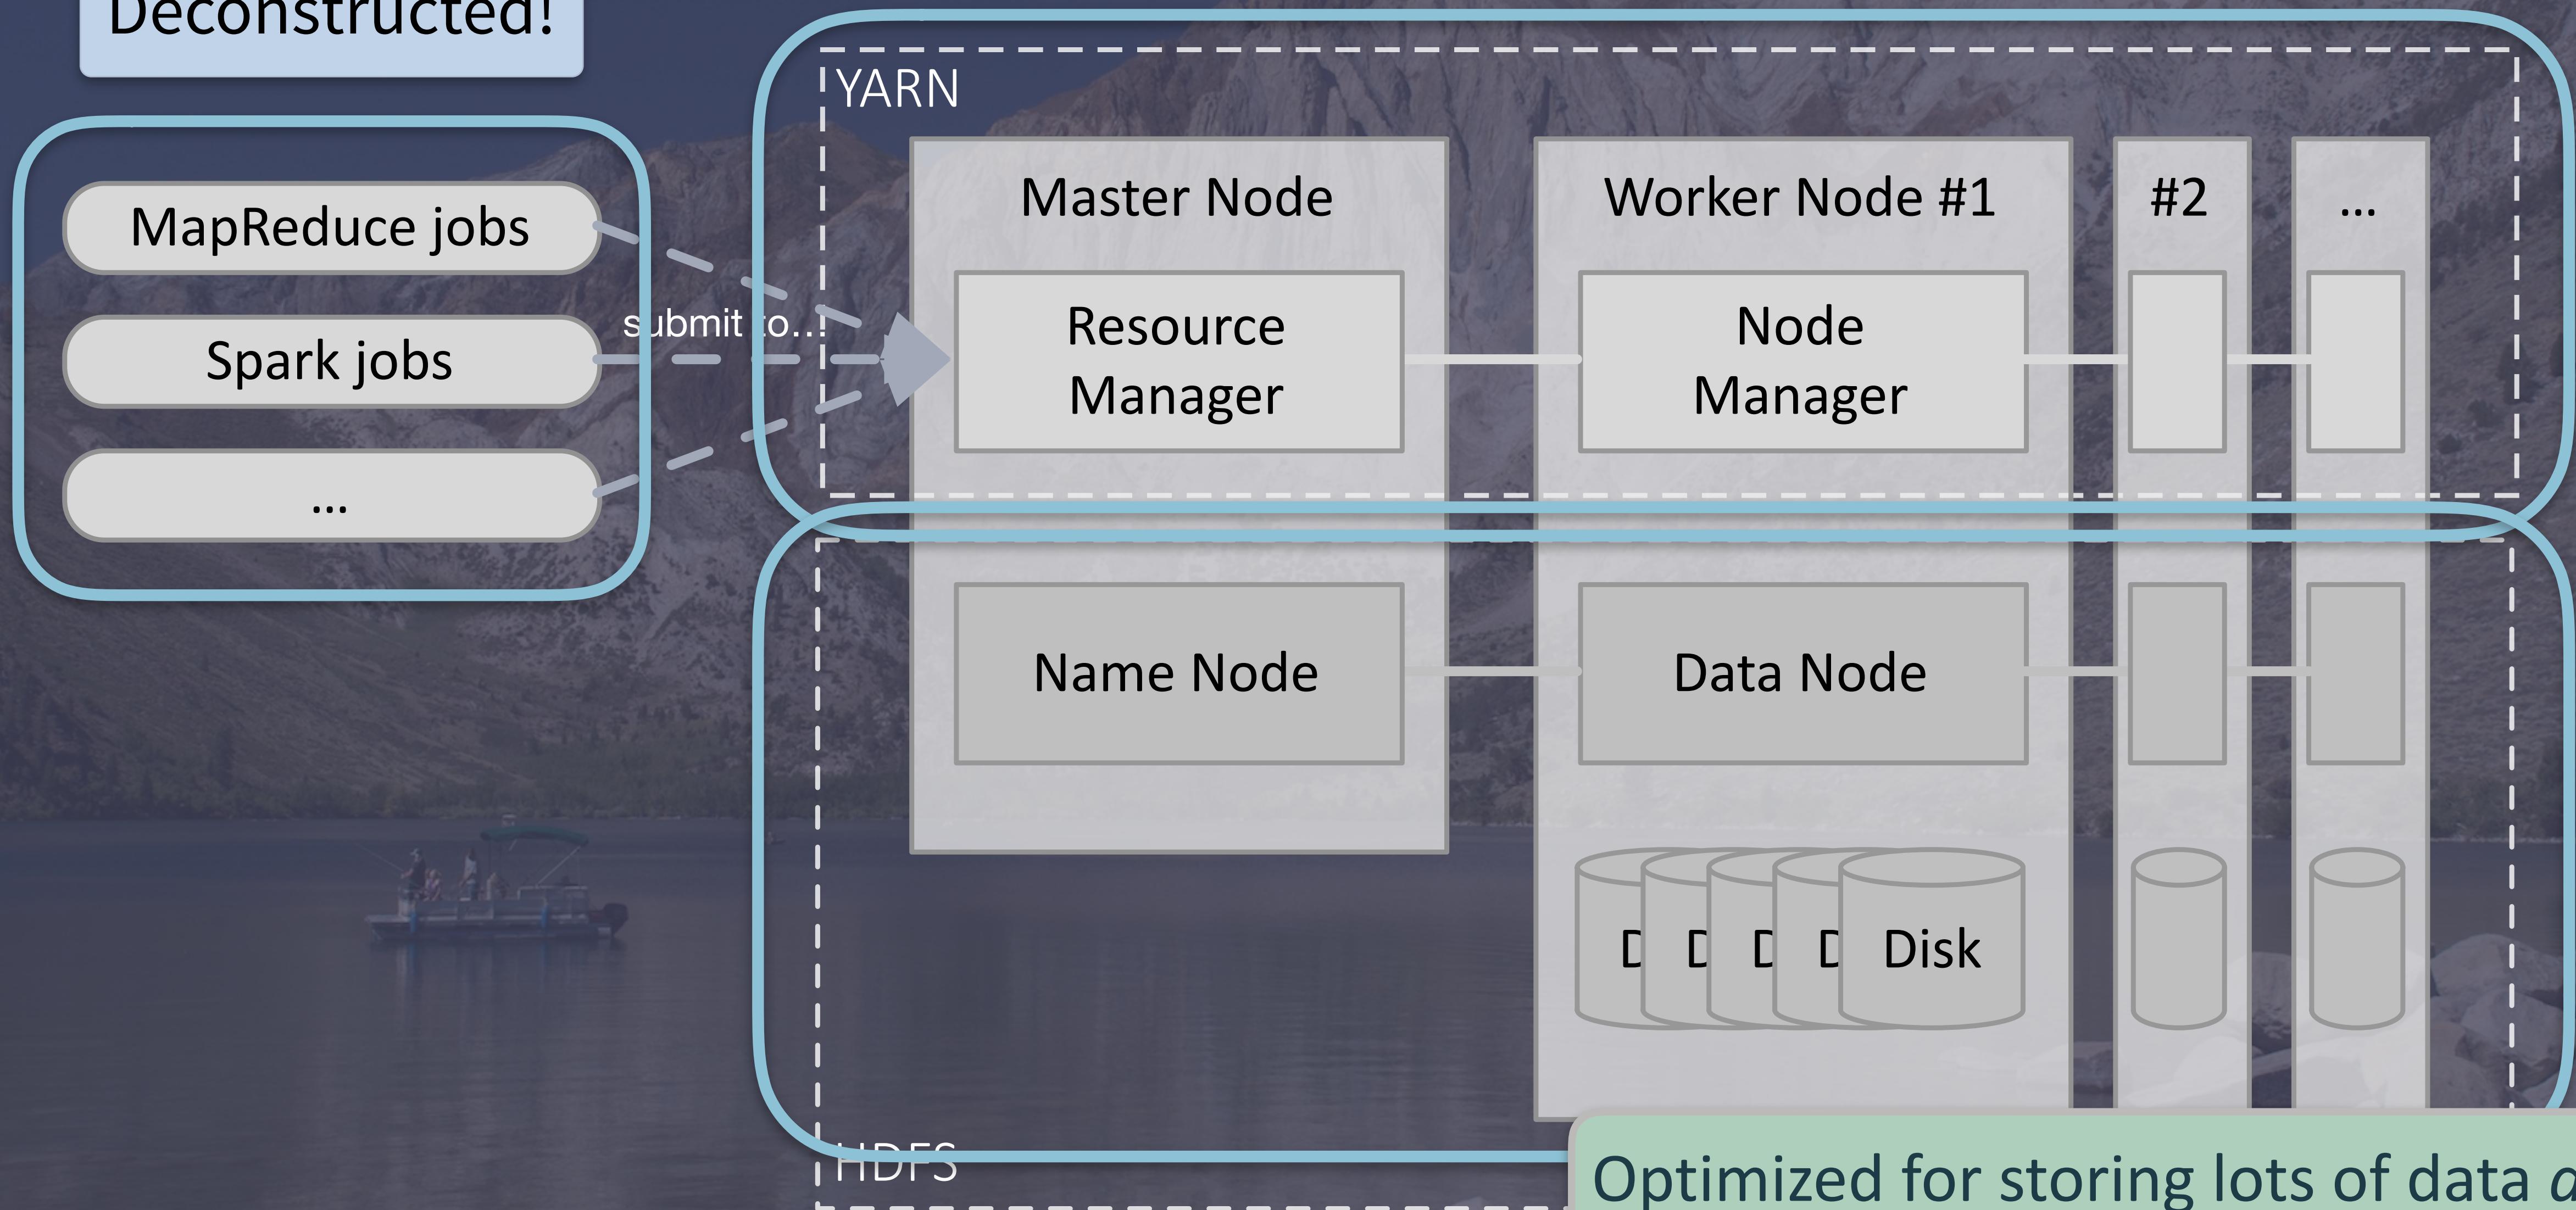
Storage

Compute



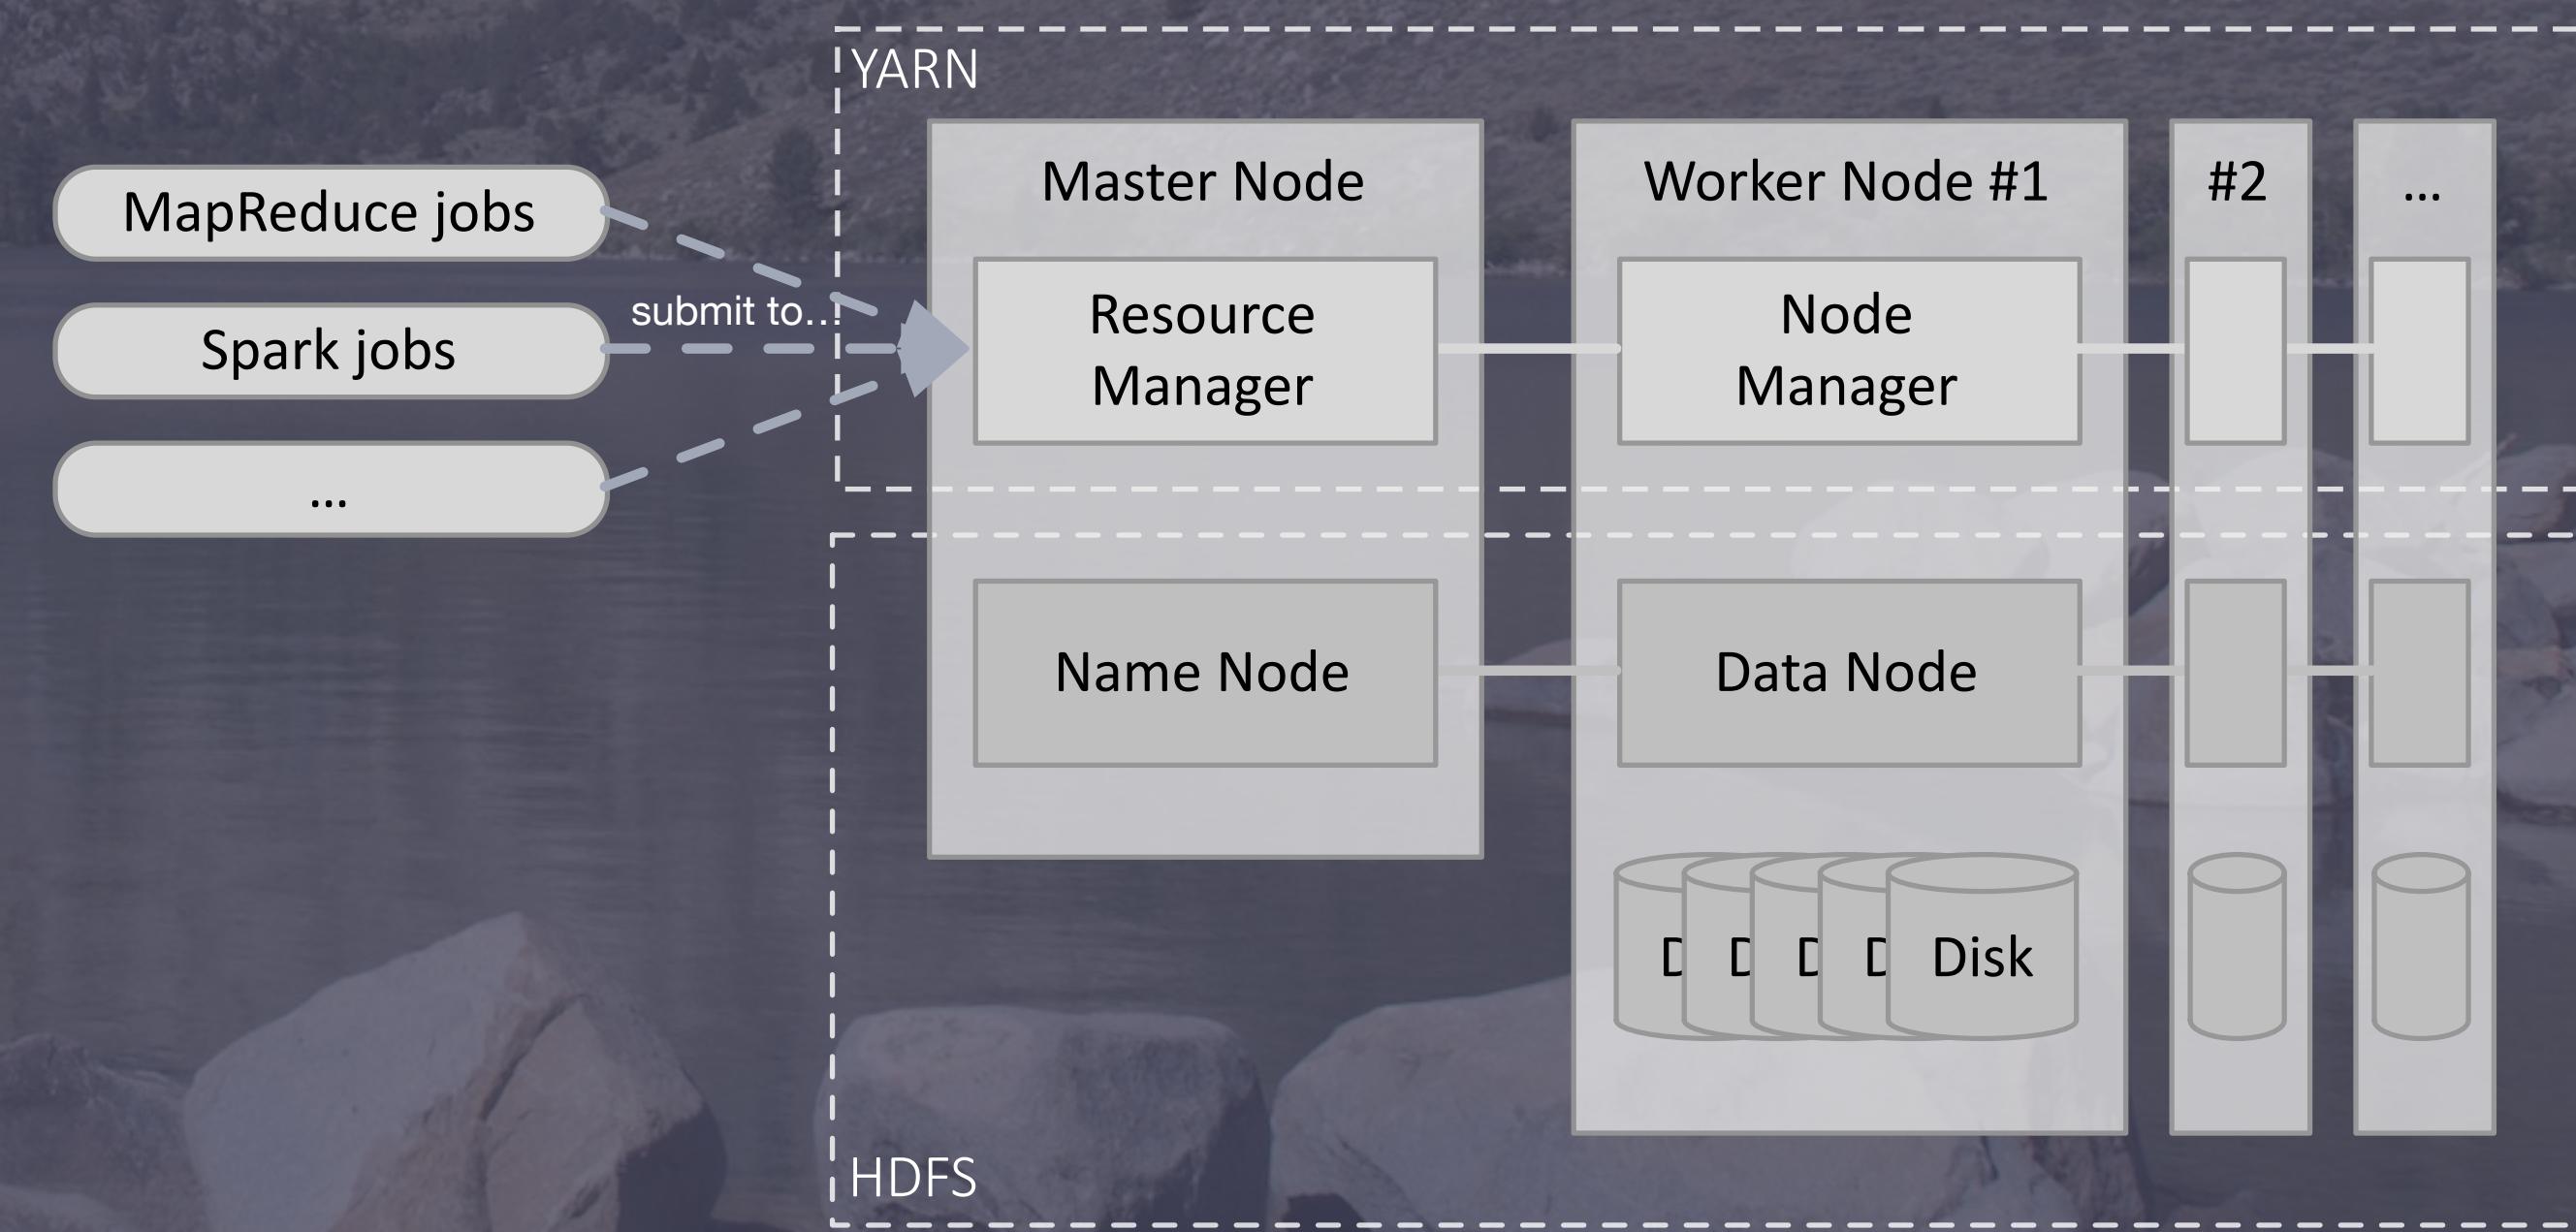


Database Deconstructed!

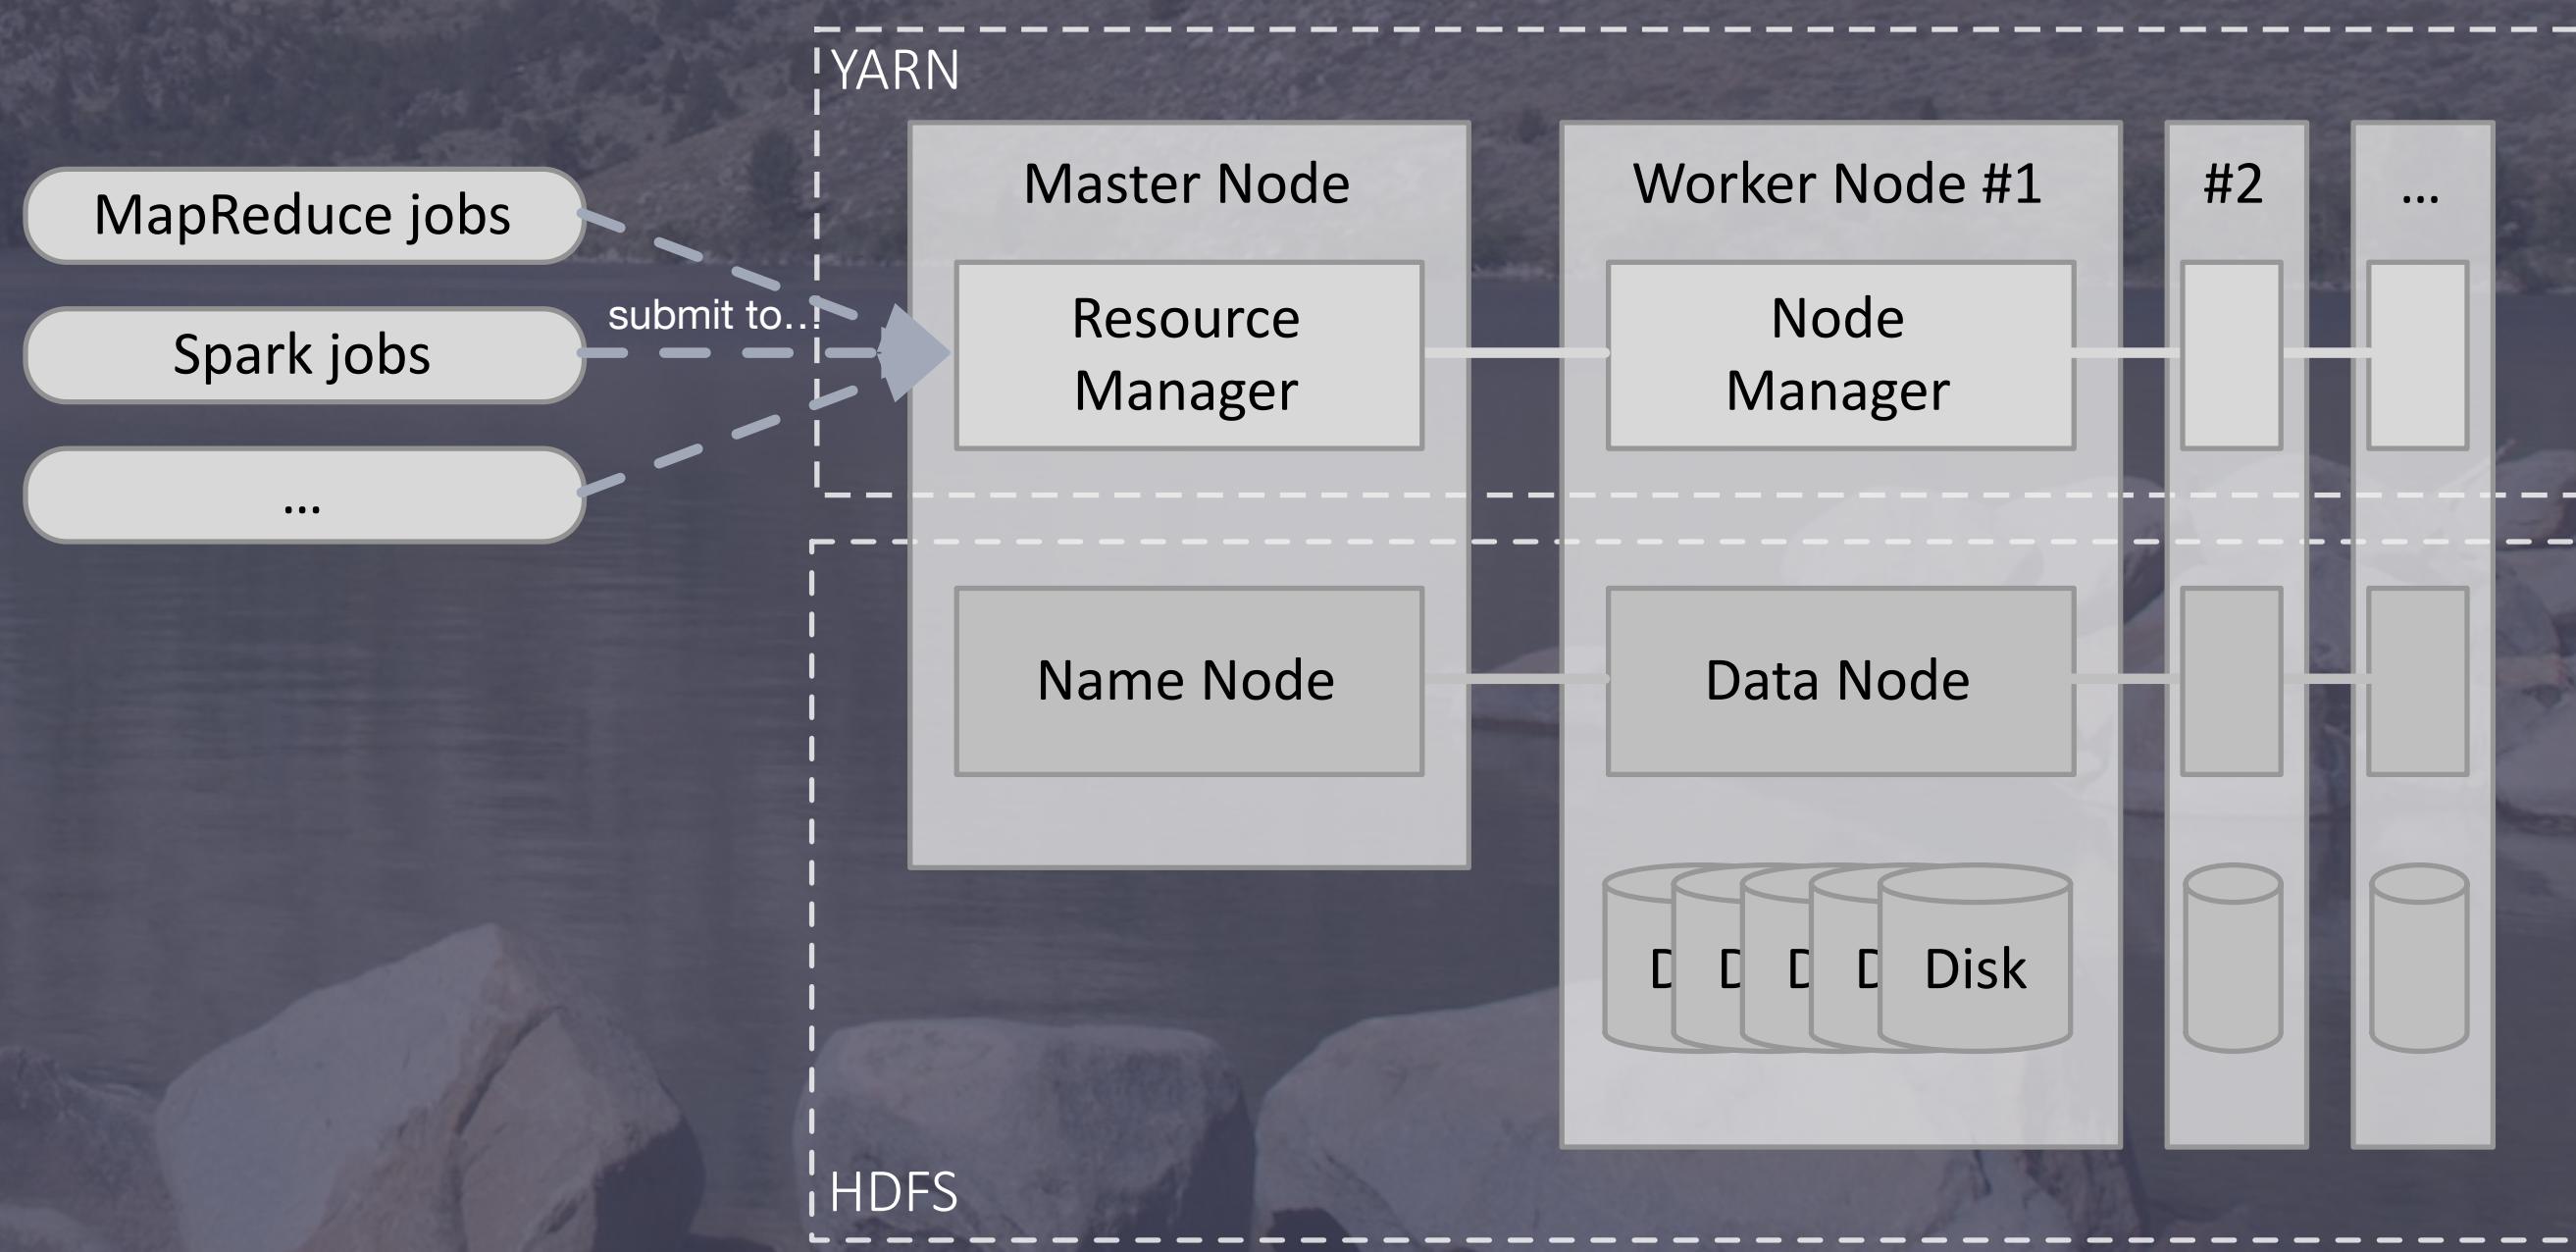


Optimized for storing lots of data *at rest*, with subsequent processing, but not optimized for data *in motion*.

- Characteristics
 - Batch and interactive
 - Massive storage - HDFS is the data “backplane”
 - Multiuser jobs

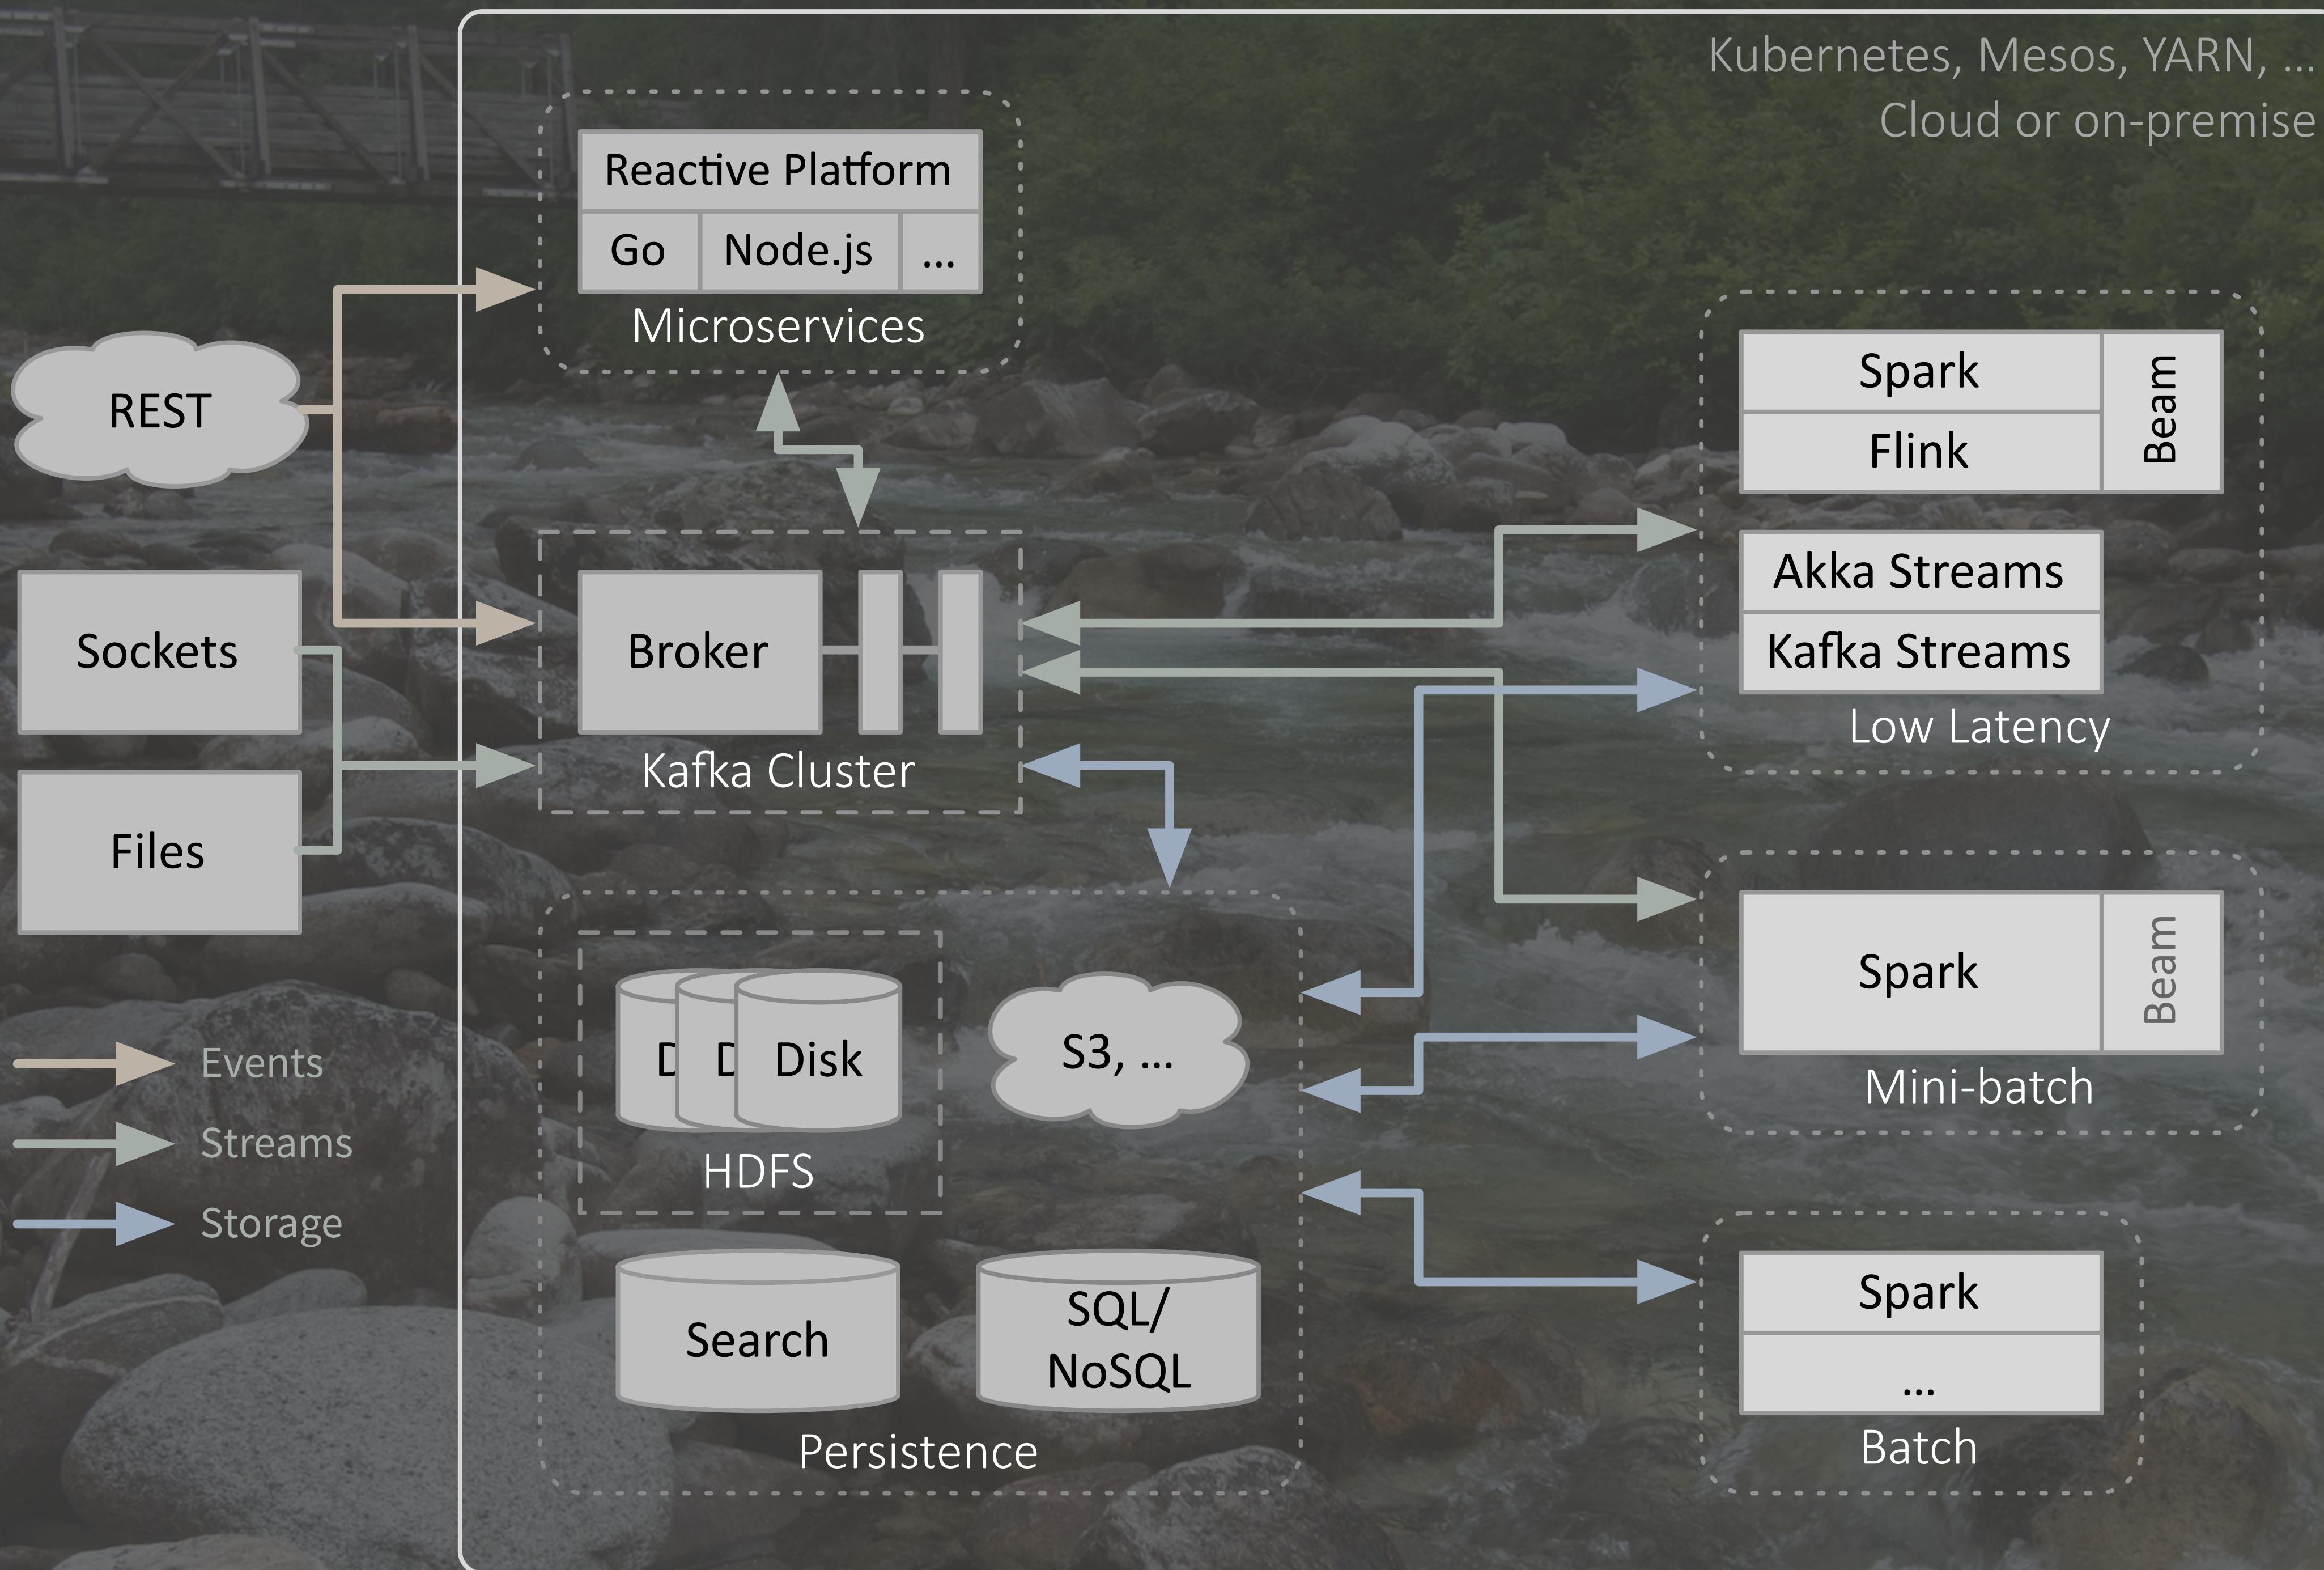


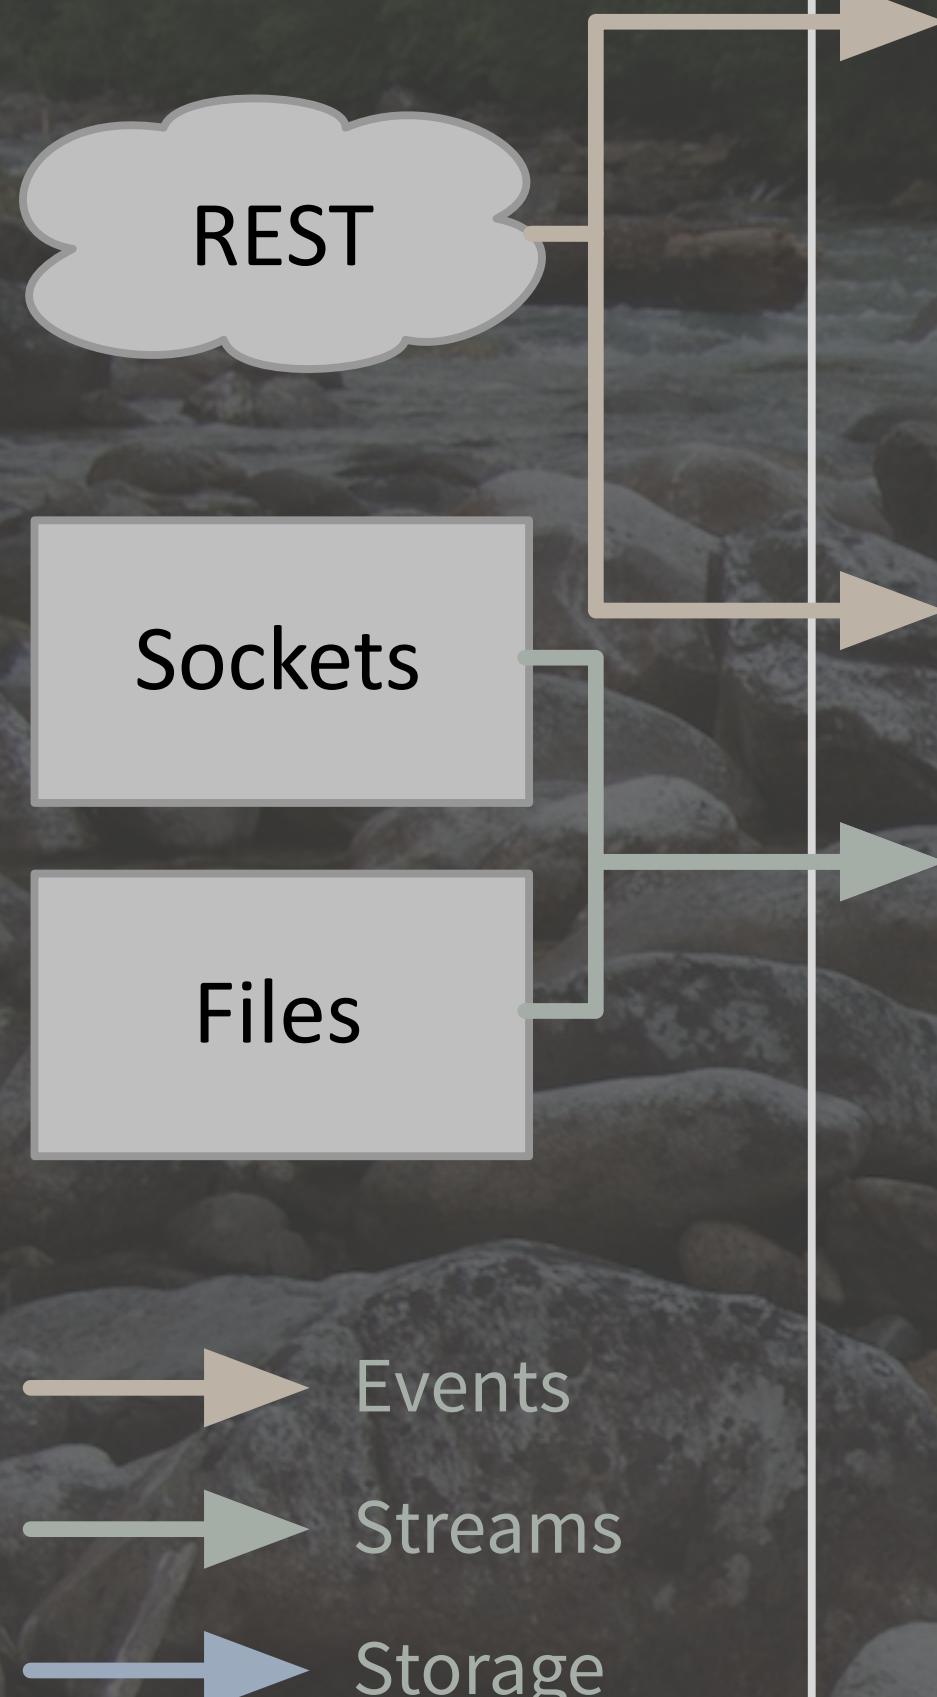
- Use Cases
 - Data warehouse replacement
 - Interactive exploration
- Offline ML model training
- ...



A scenic view of a rocky river flowing through a lush green forest. The river is filled with large, smooth stones and rocks, with white water cascading over them. In the background, a rustic wooden bridge spans the river. The surrounding trees are dense and vibrant green.

New Streaming, “Fast Data” Architecture





Kubernetes and Mesos provide the job and resource management needed for dynamic, heterogenous work loads

Broker
Kafka Cluster

Disk

HDFS

Search

S3, ...

SQL/
NoSQL

Persistence

Kubernetes, Mesos, YARN, ...
Cloud or on-premise

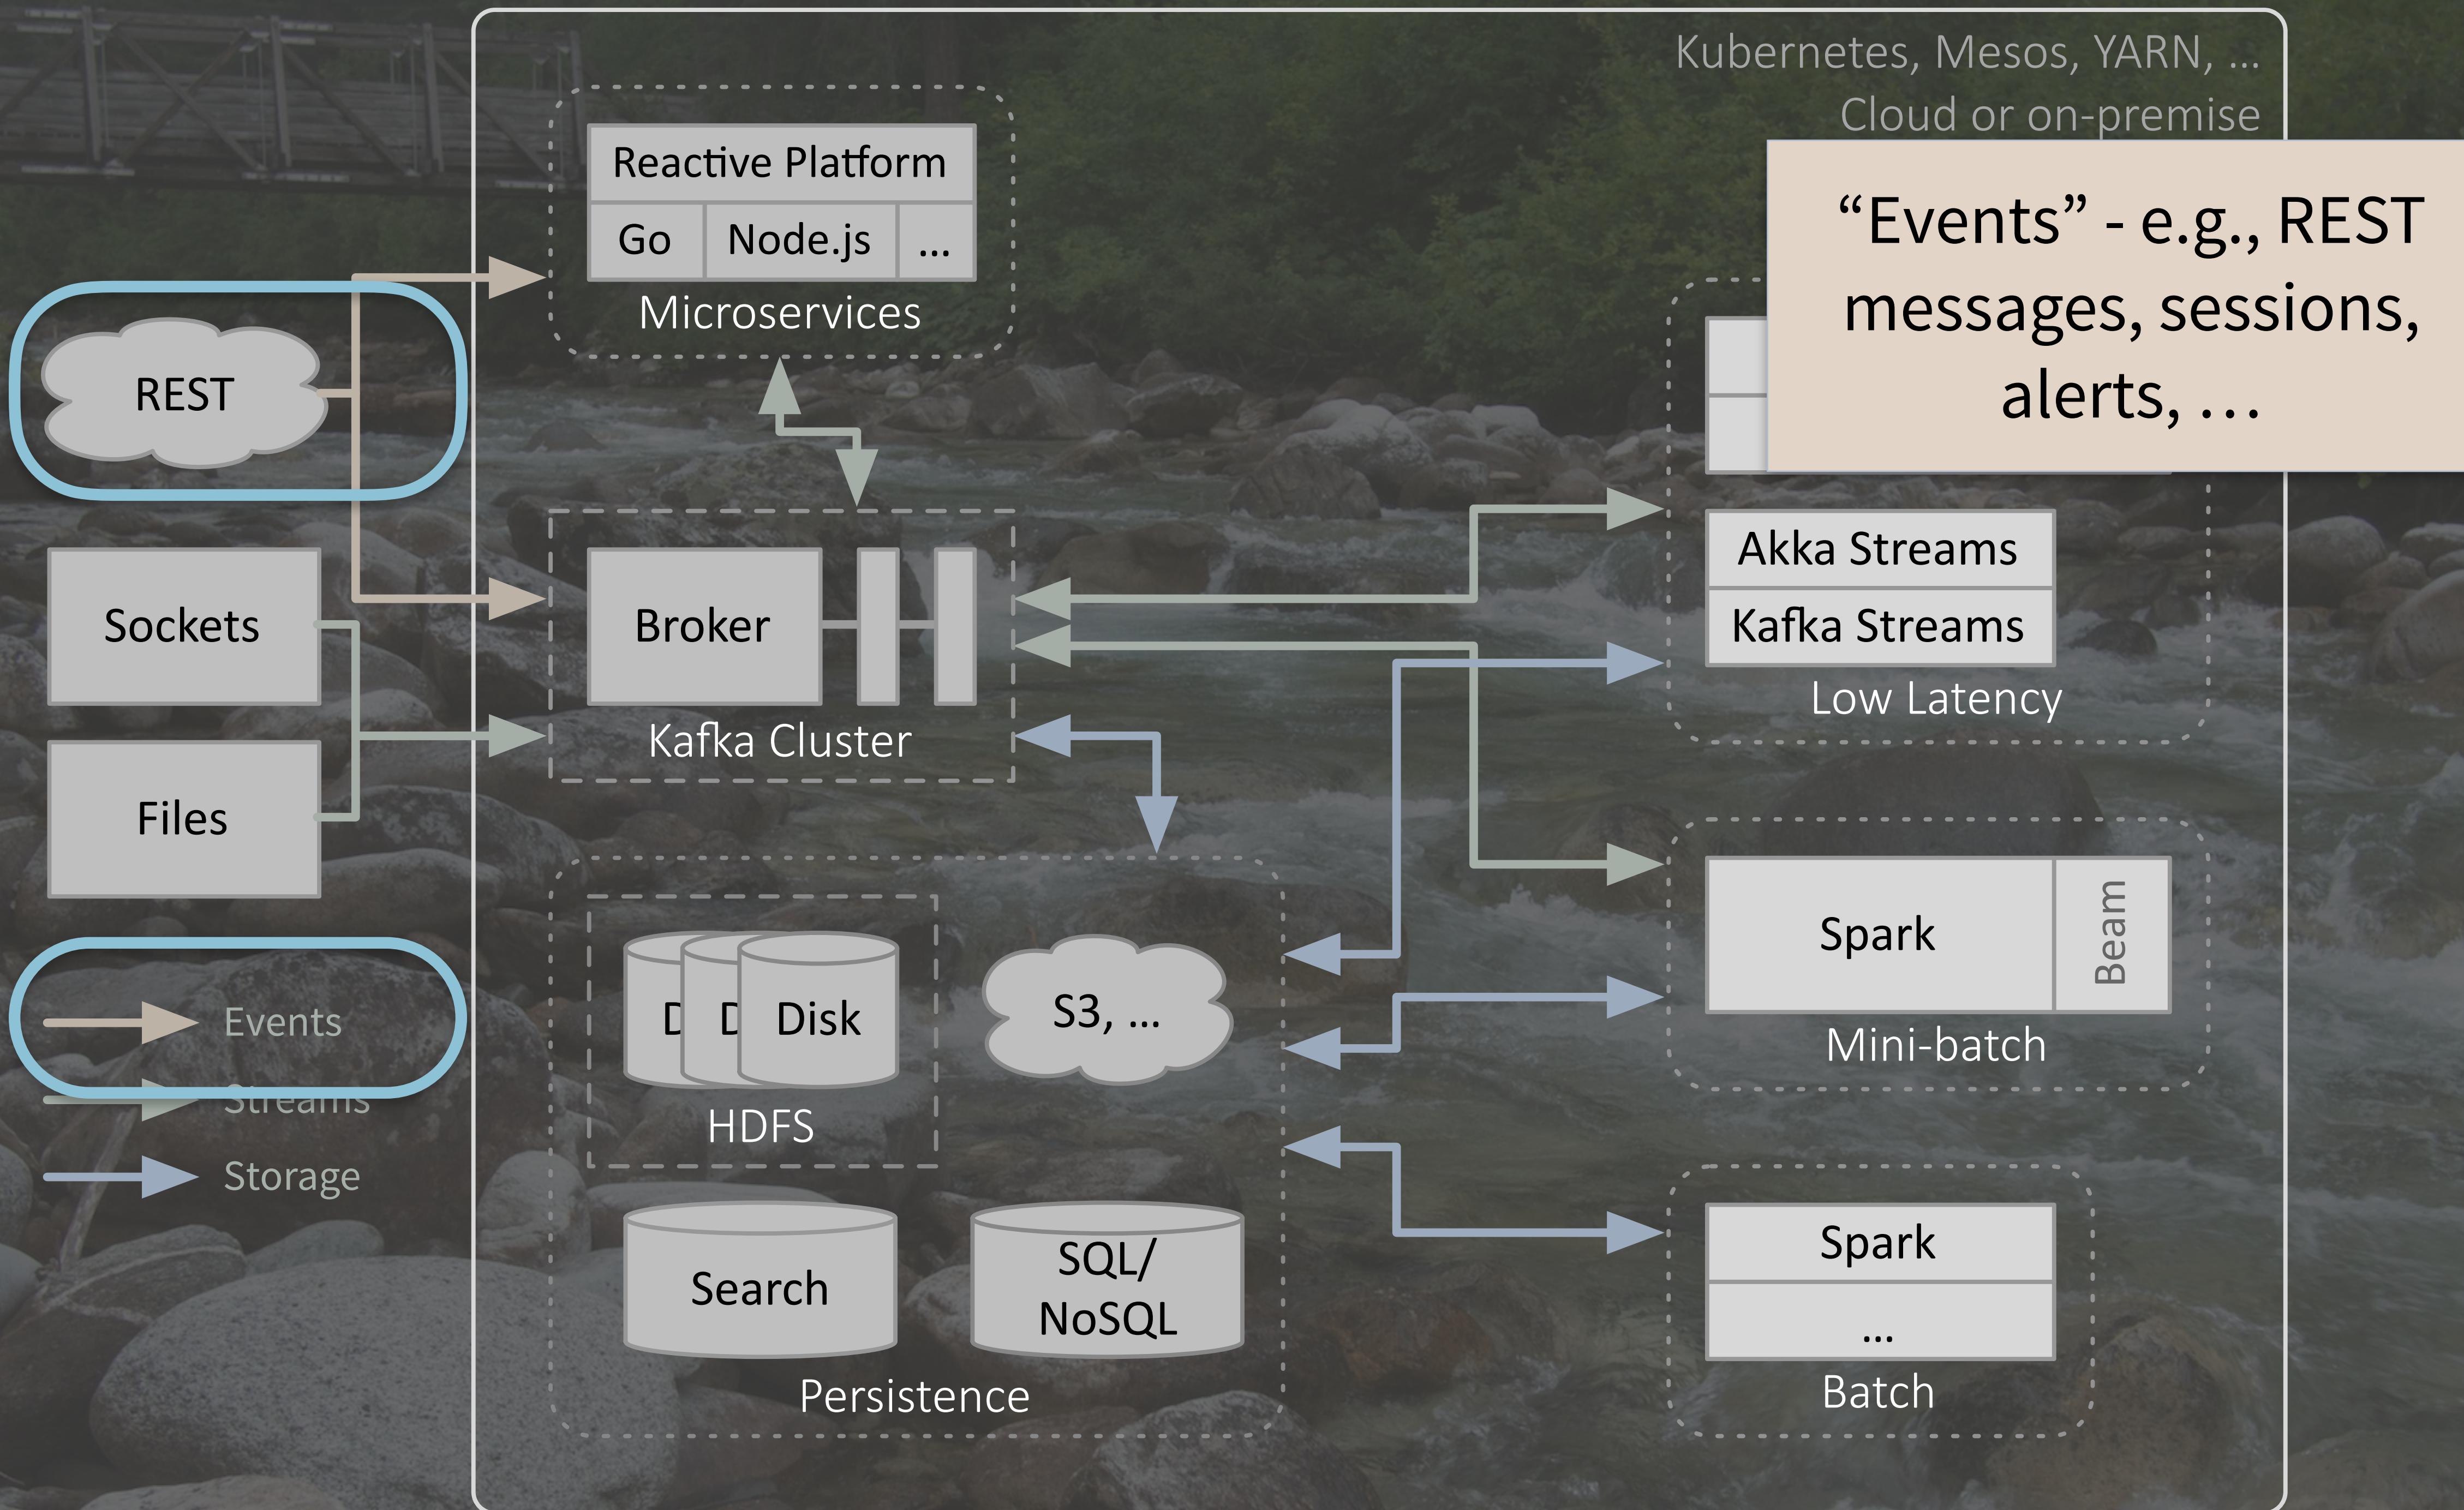
While YARN can be used, it's not flexible enough for today's dynamic workloads

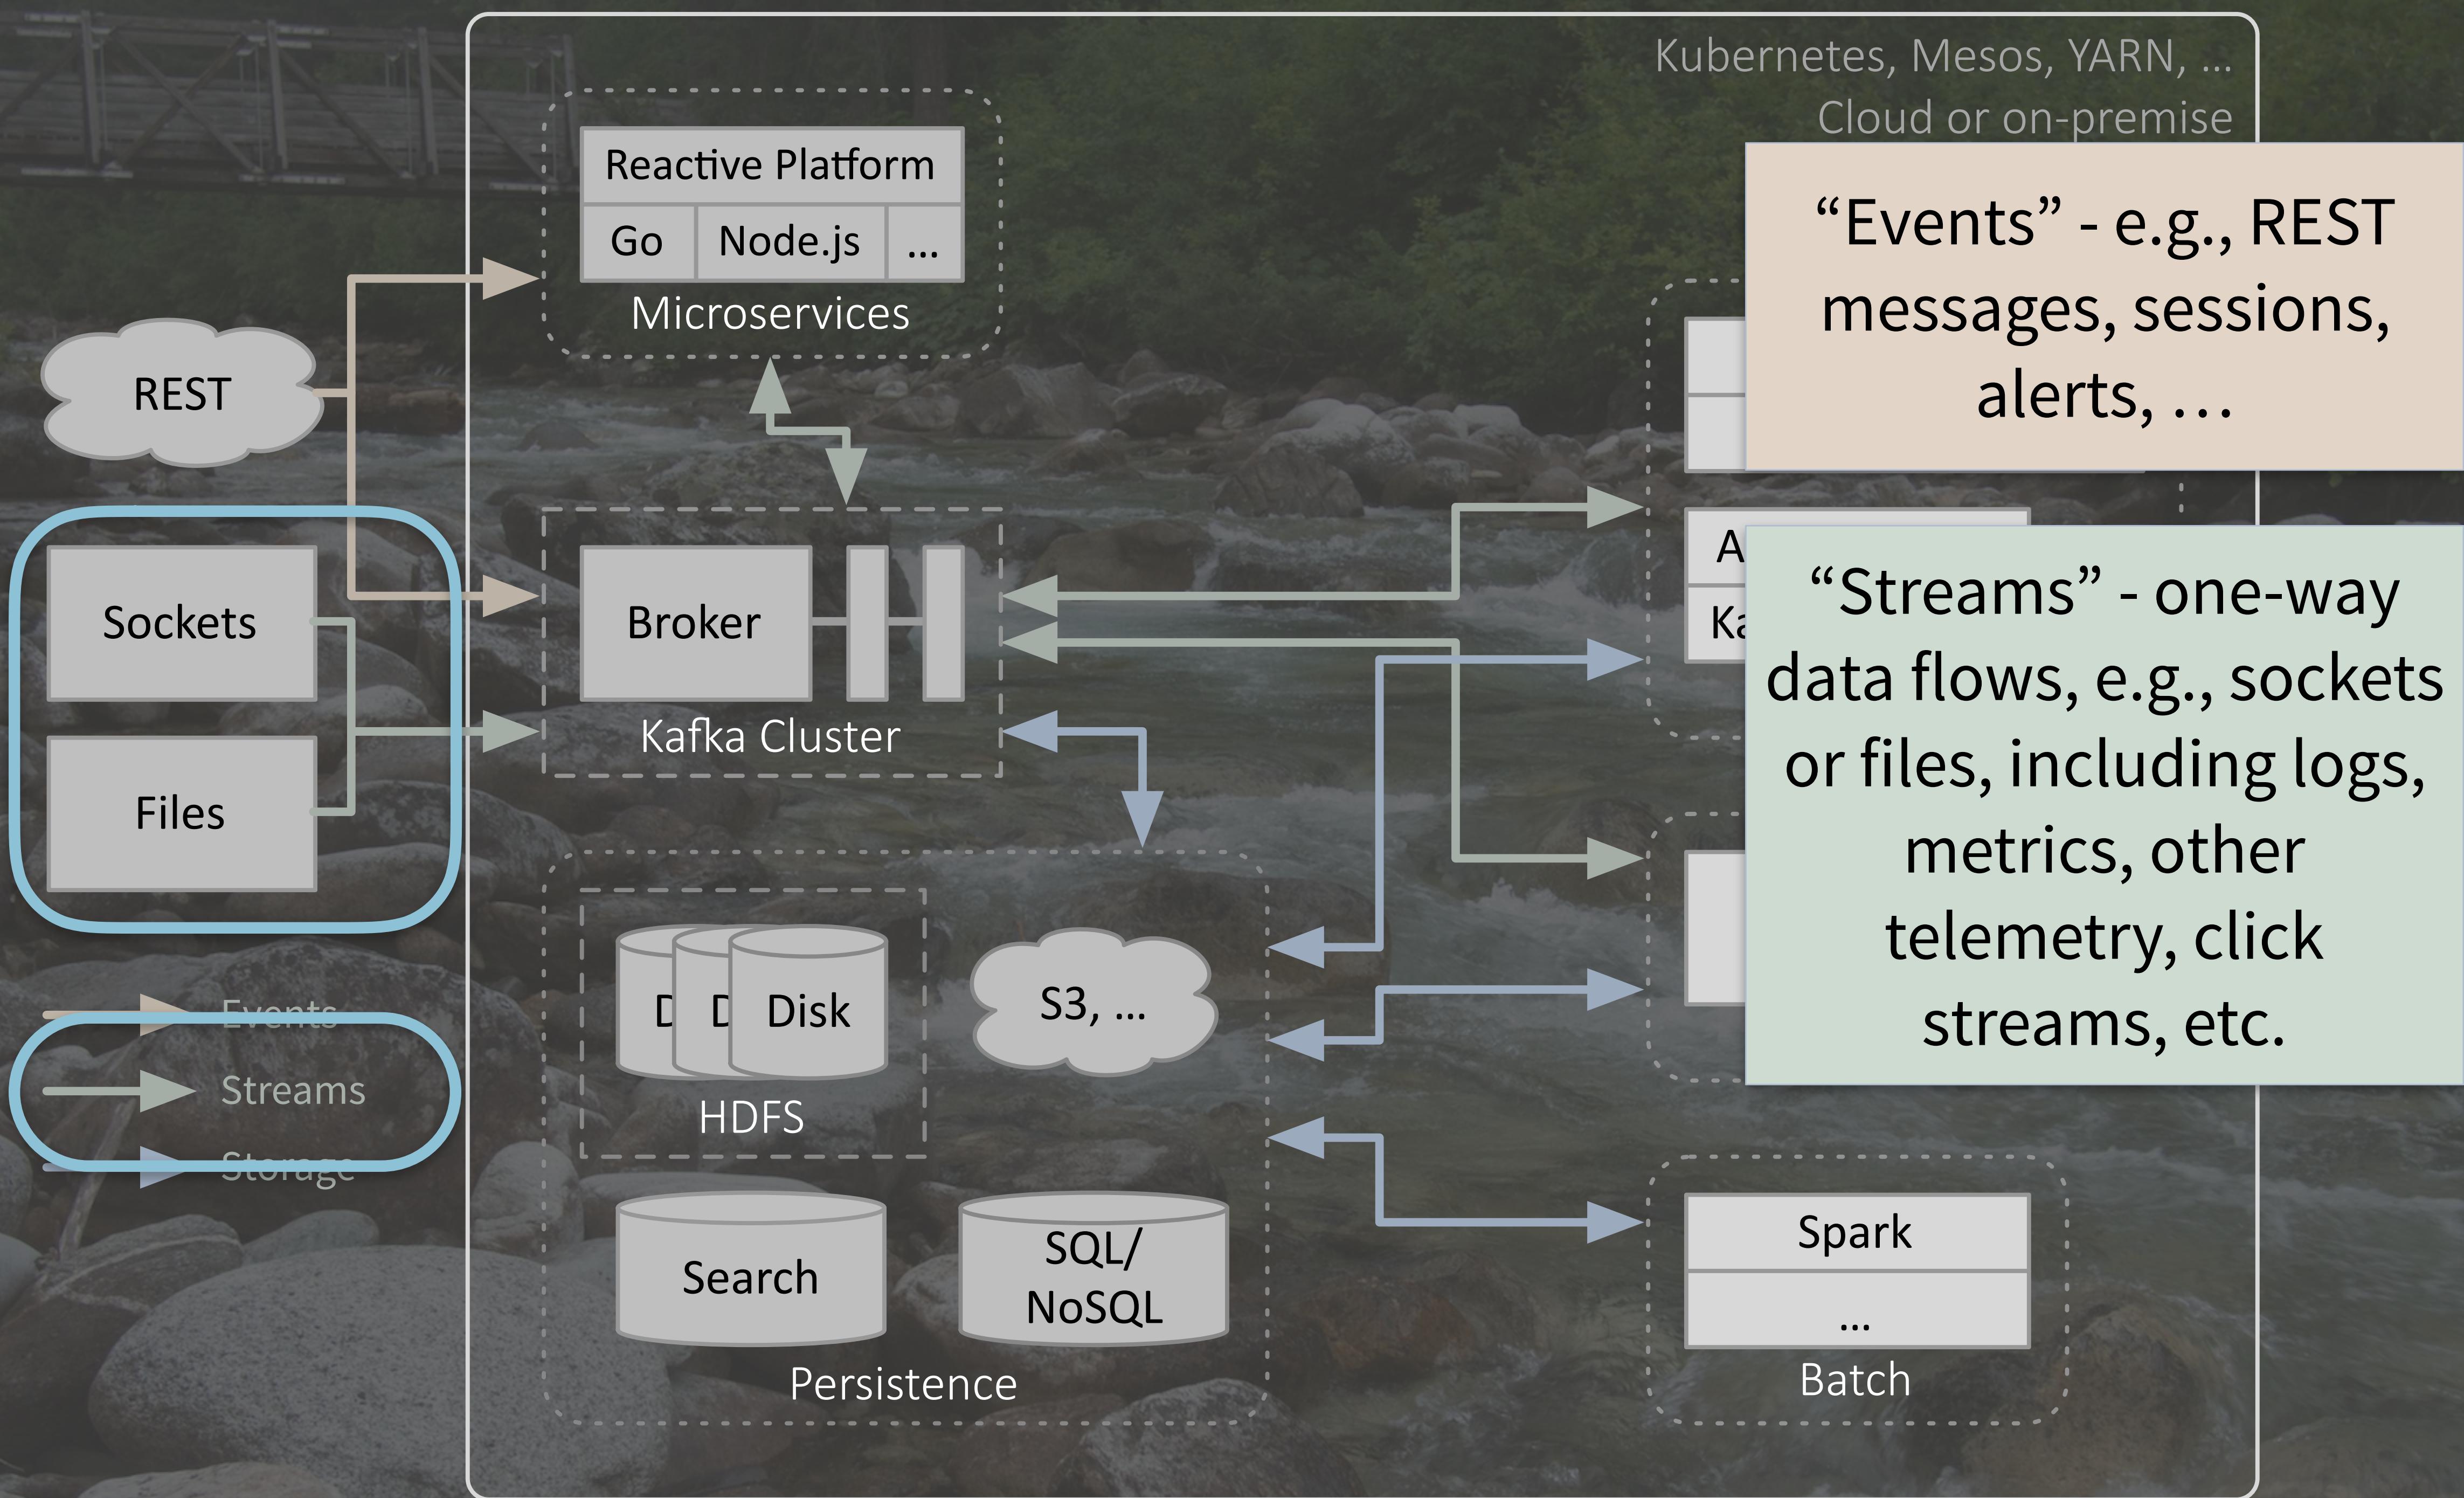
Deploy in the cloud or on premise

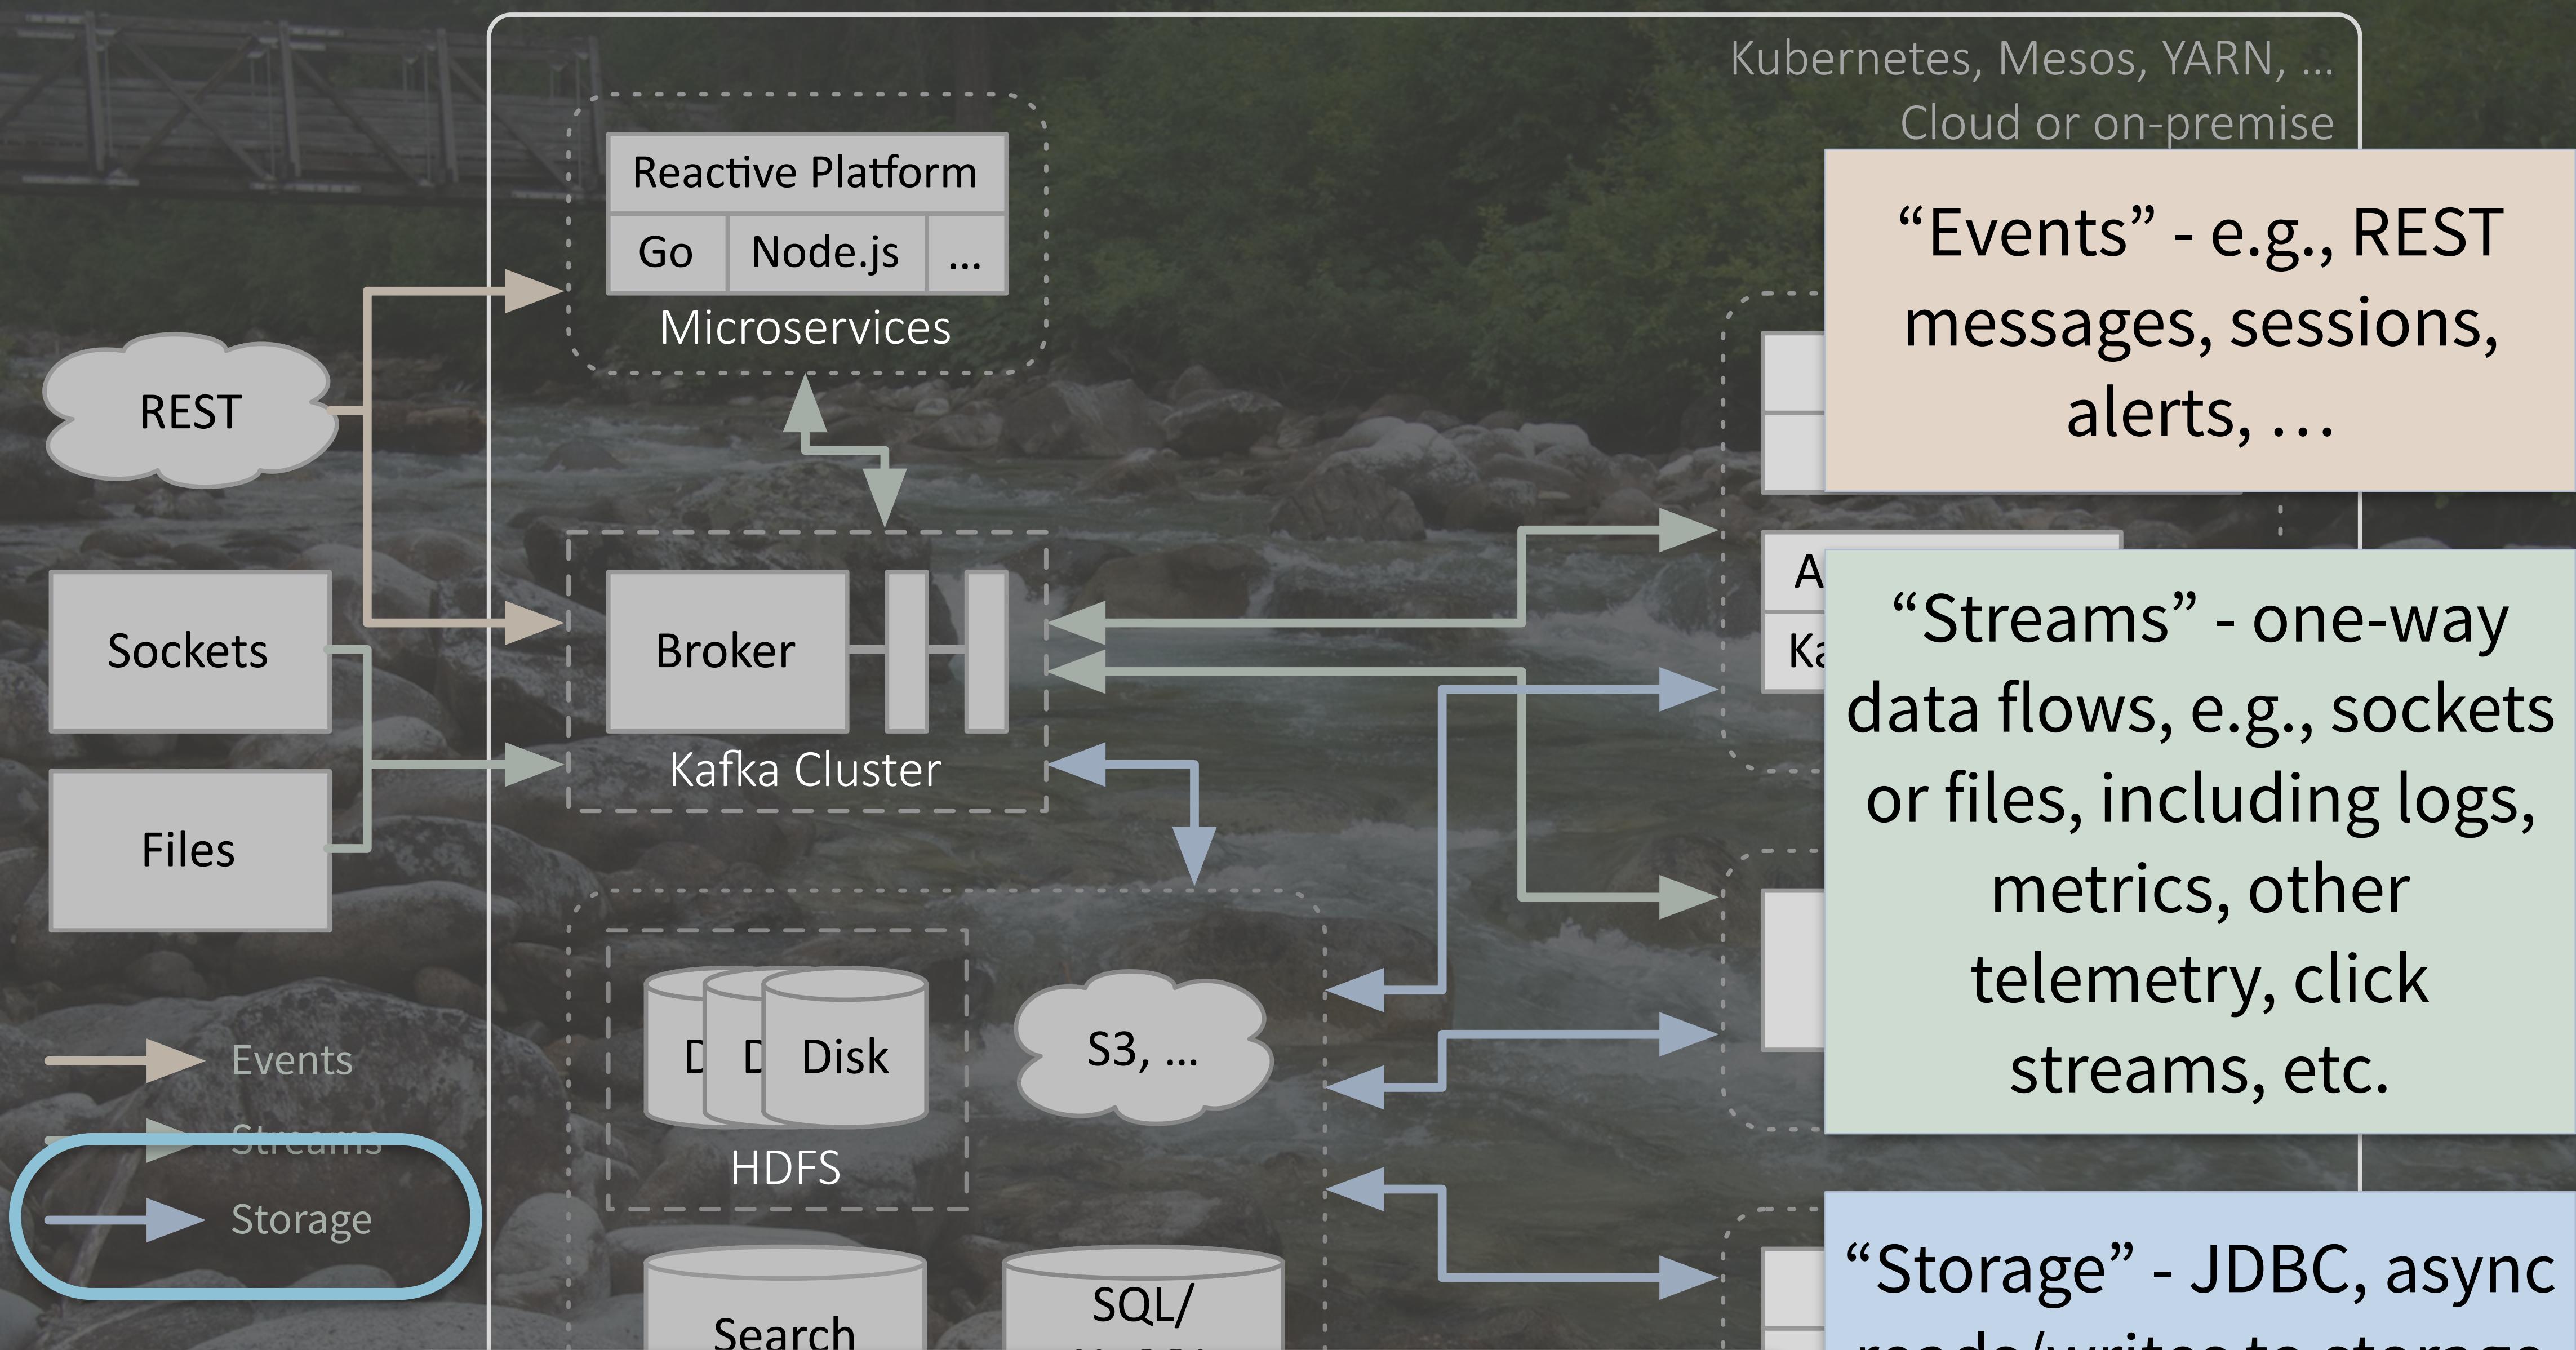
Spark
Mini-batch

Beam

Spark
...
Batch







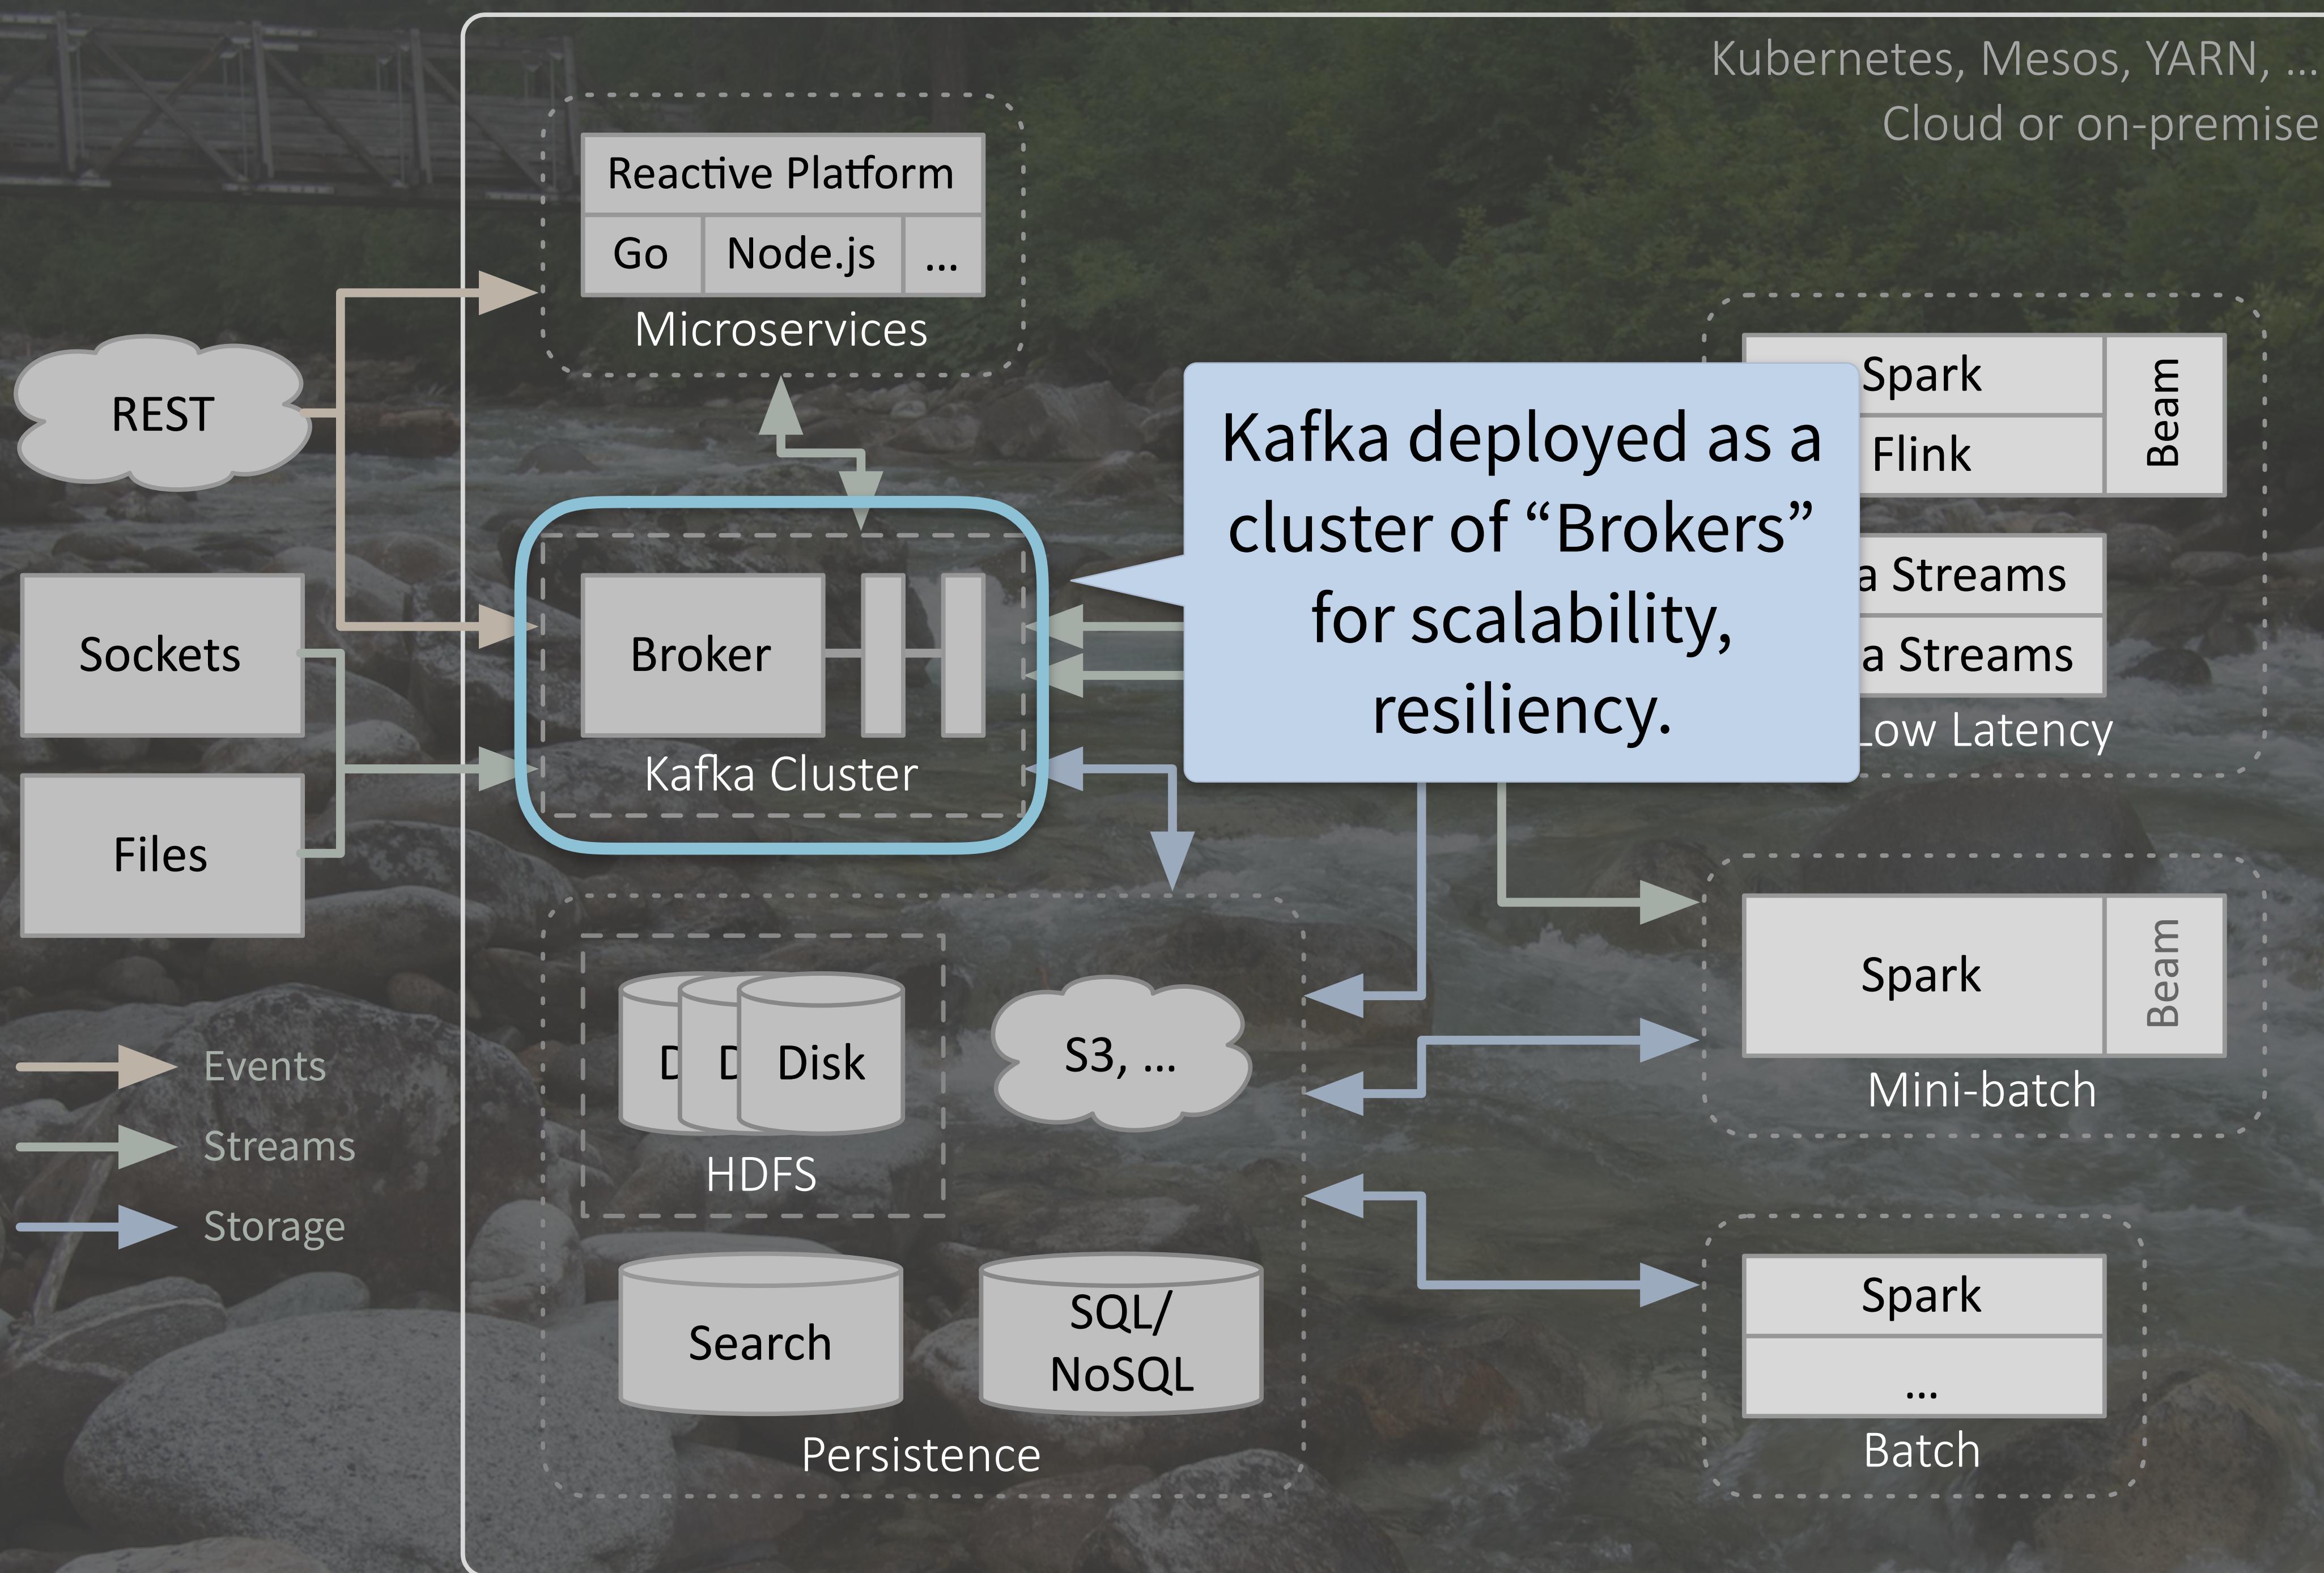
Kubernetes, Mesos, YARN, ...
Cloud or on-premise

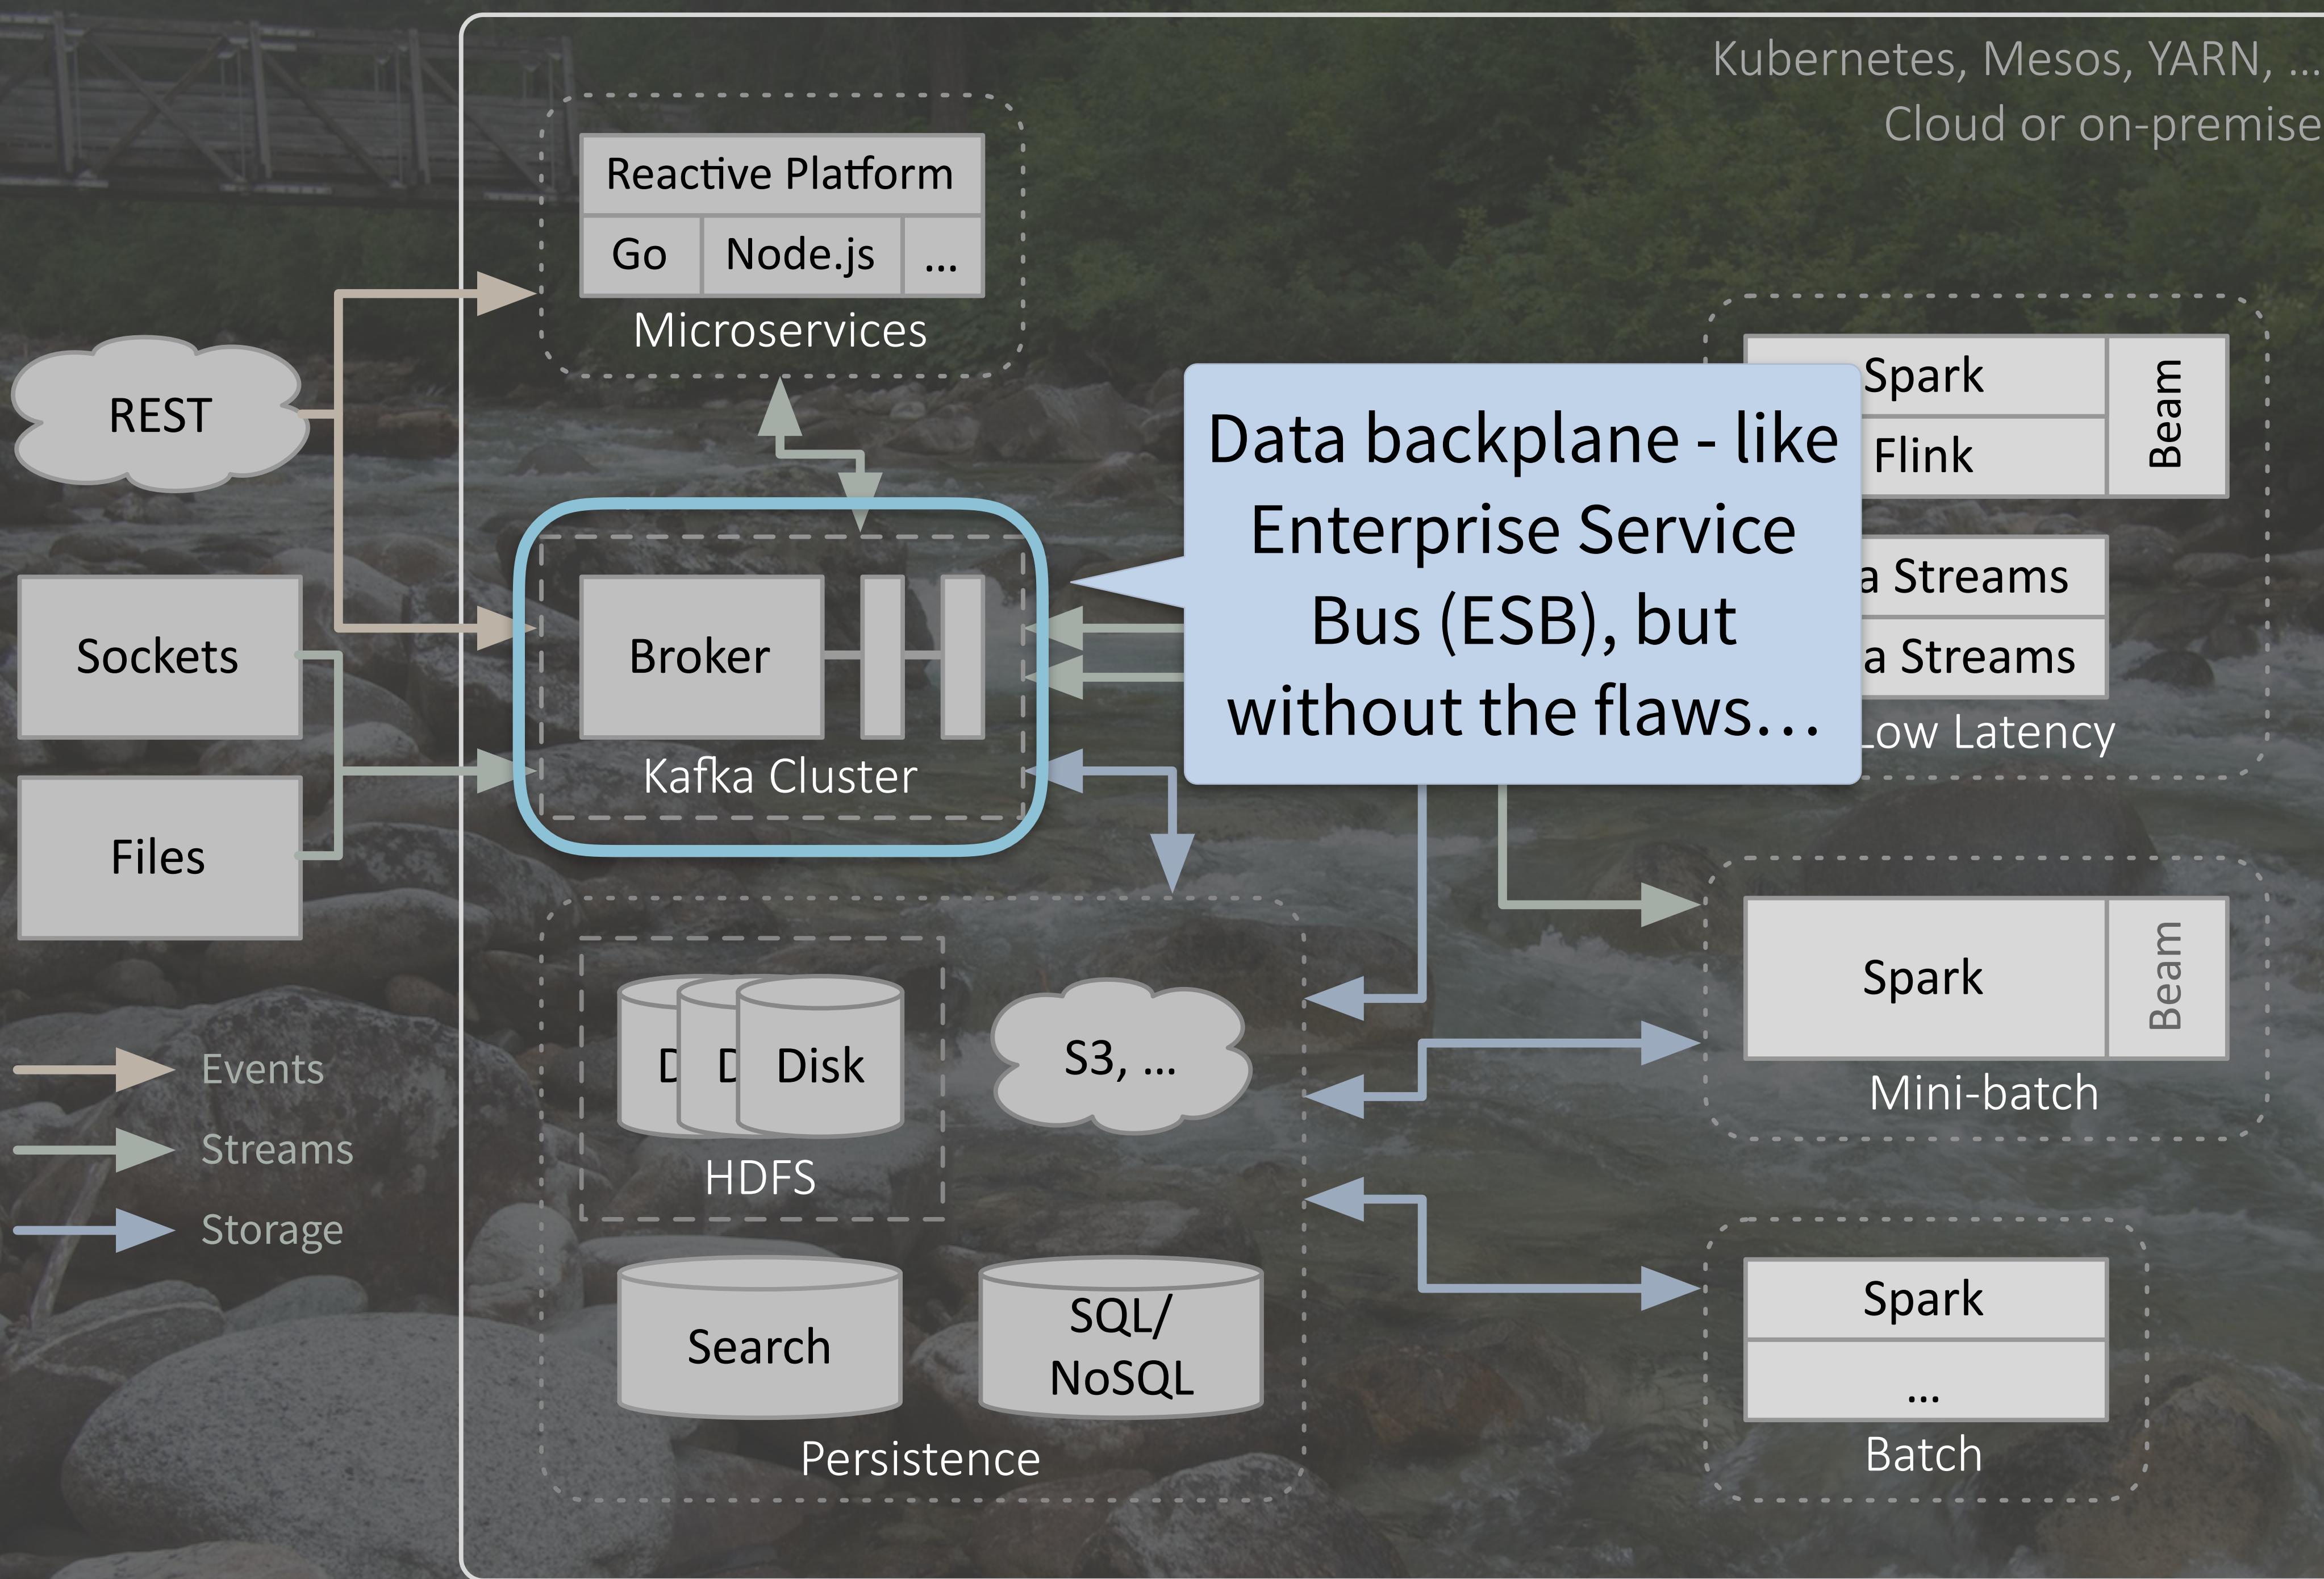
“Events” - e.g., REST messages, sessions, alerts, ...

“Streams” - one-way data flows, e.g., sockets or files, including logs, metrics, other telemetry, click streams, etc.

“Storage” - JDBC, async reads/writes to storage

Each has different volumes, velocities, latency characteristics, protocols, etc.

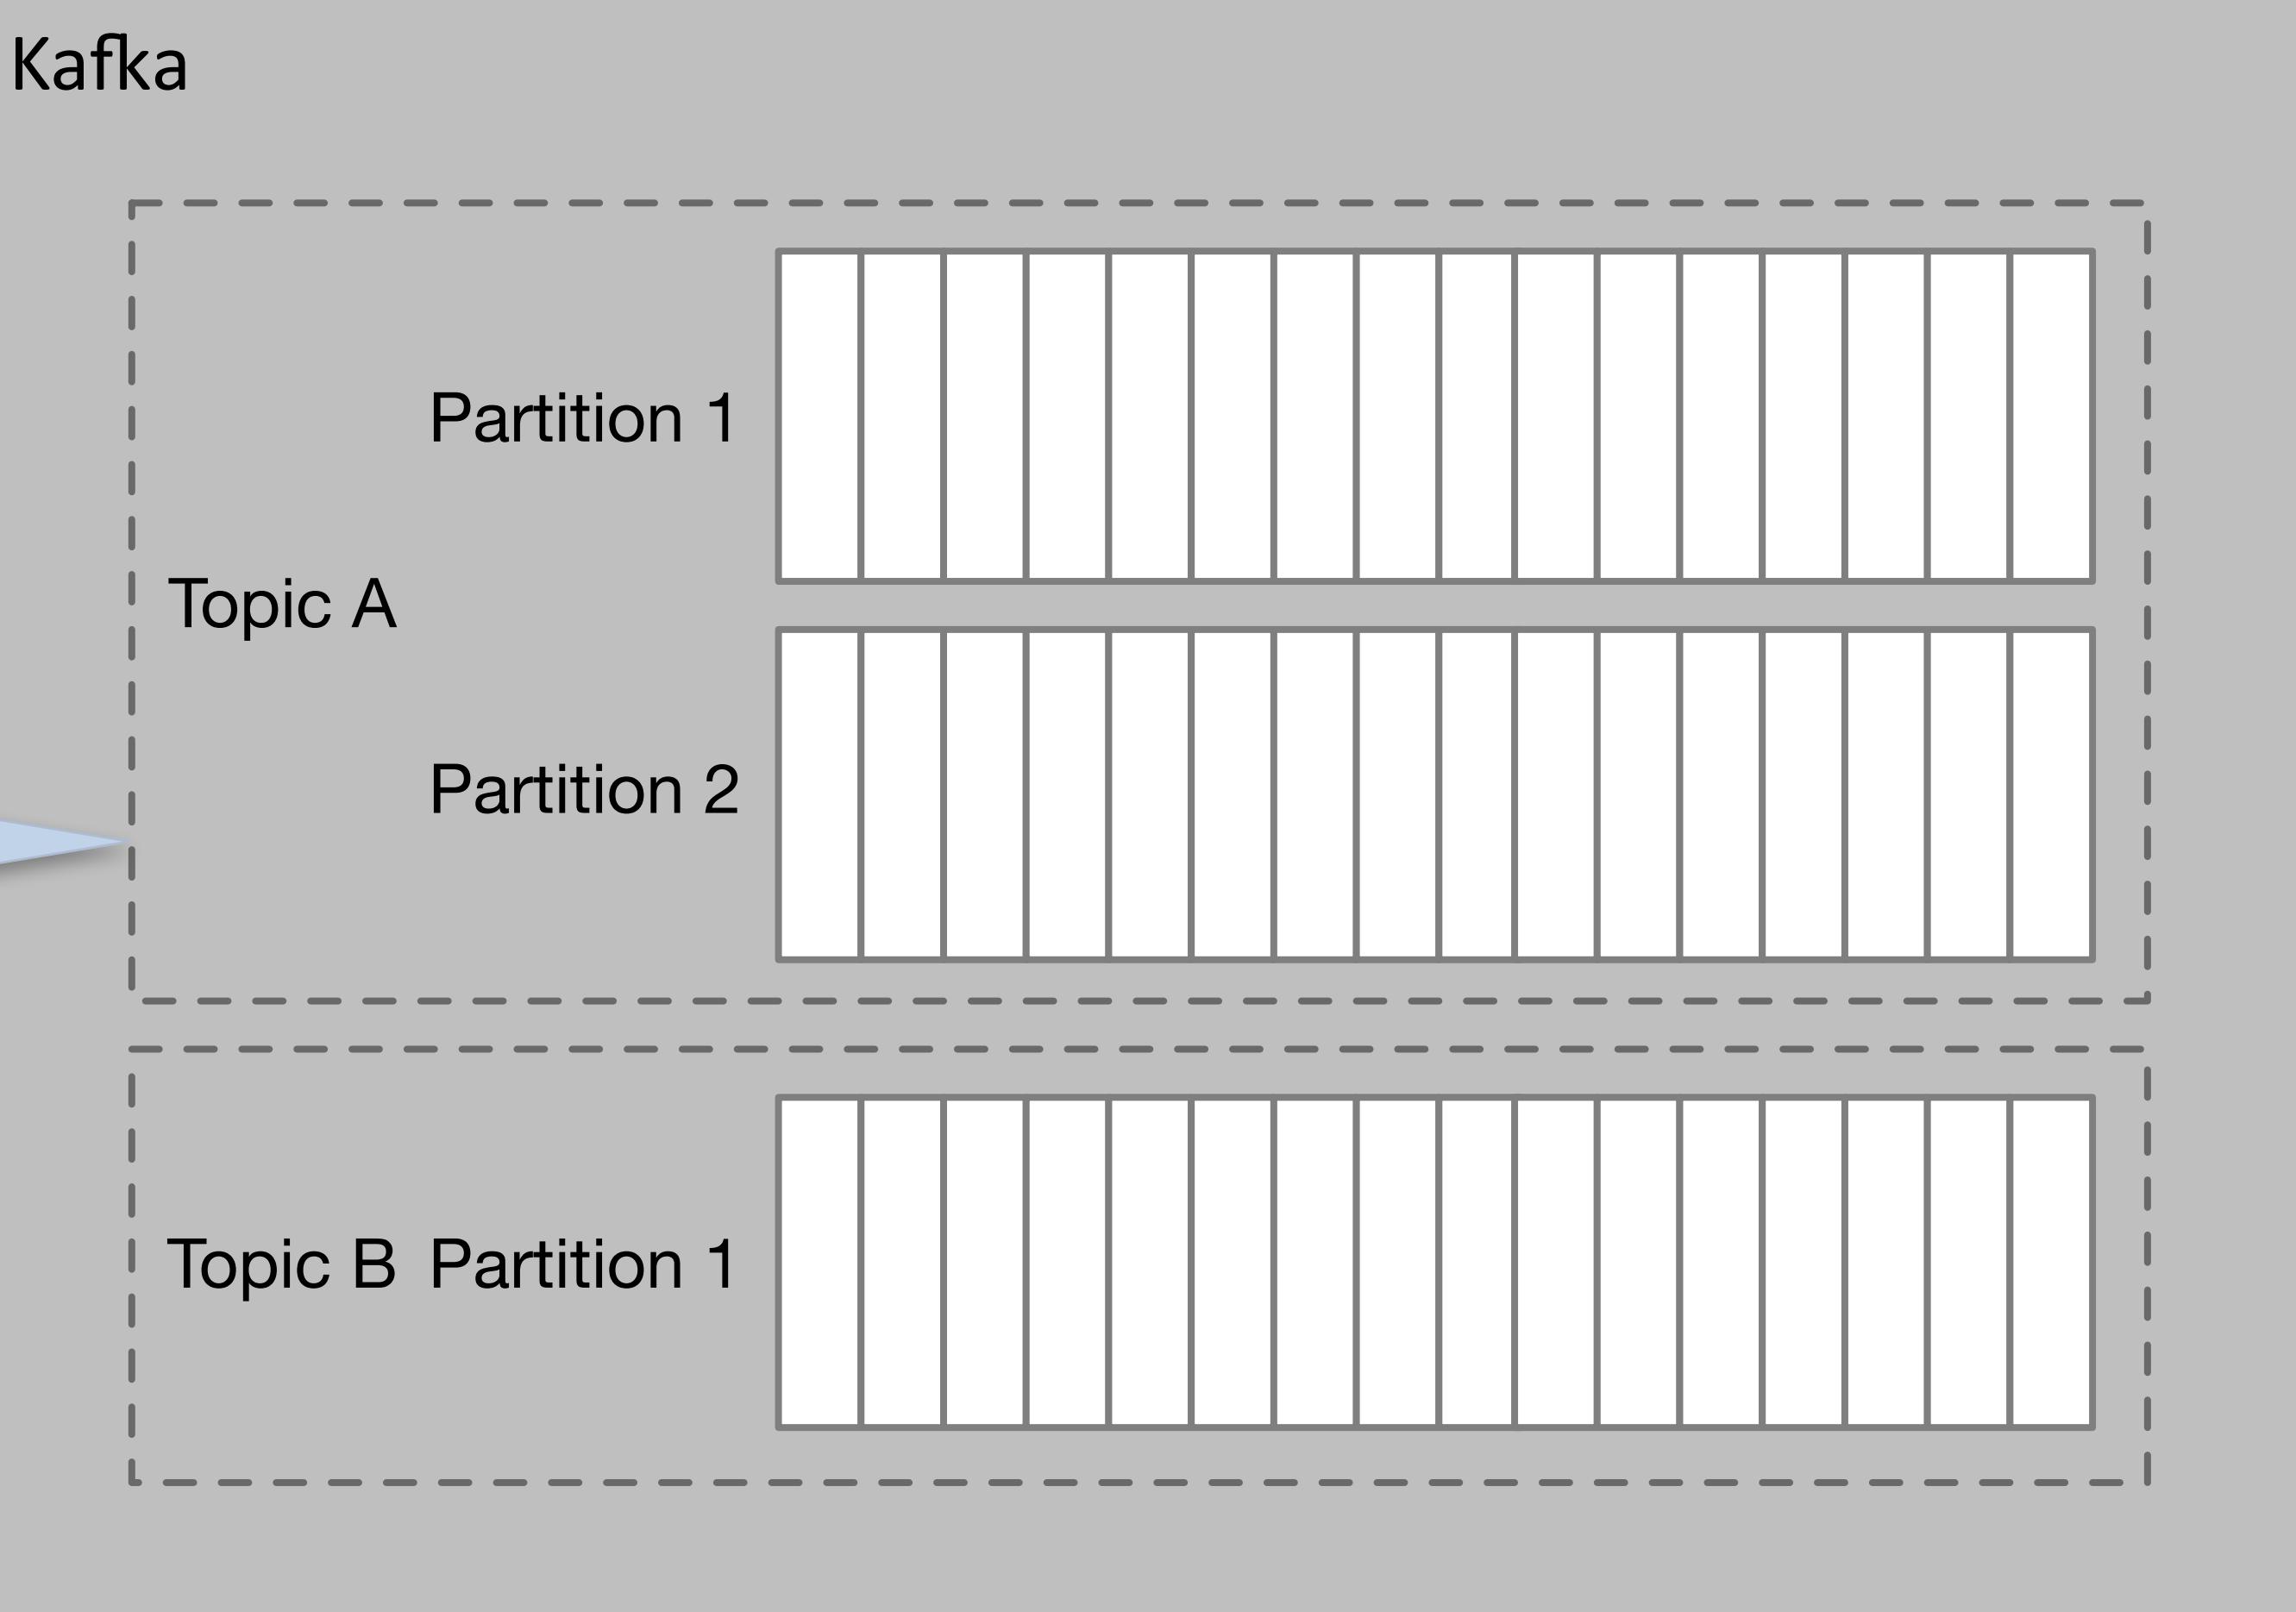




Why Kafka?

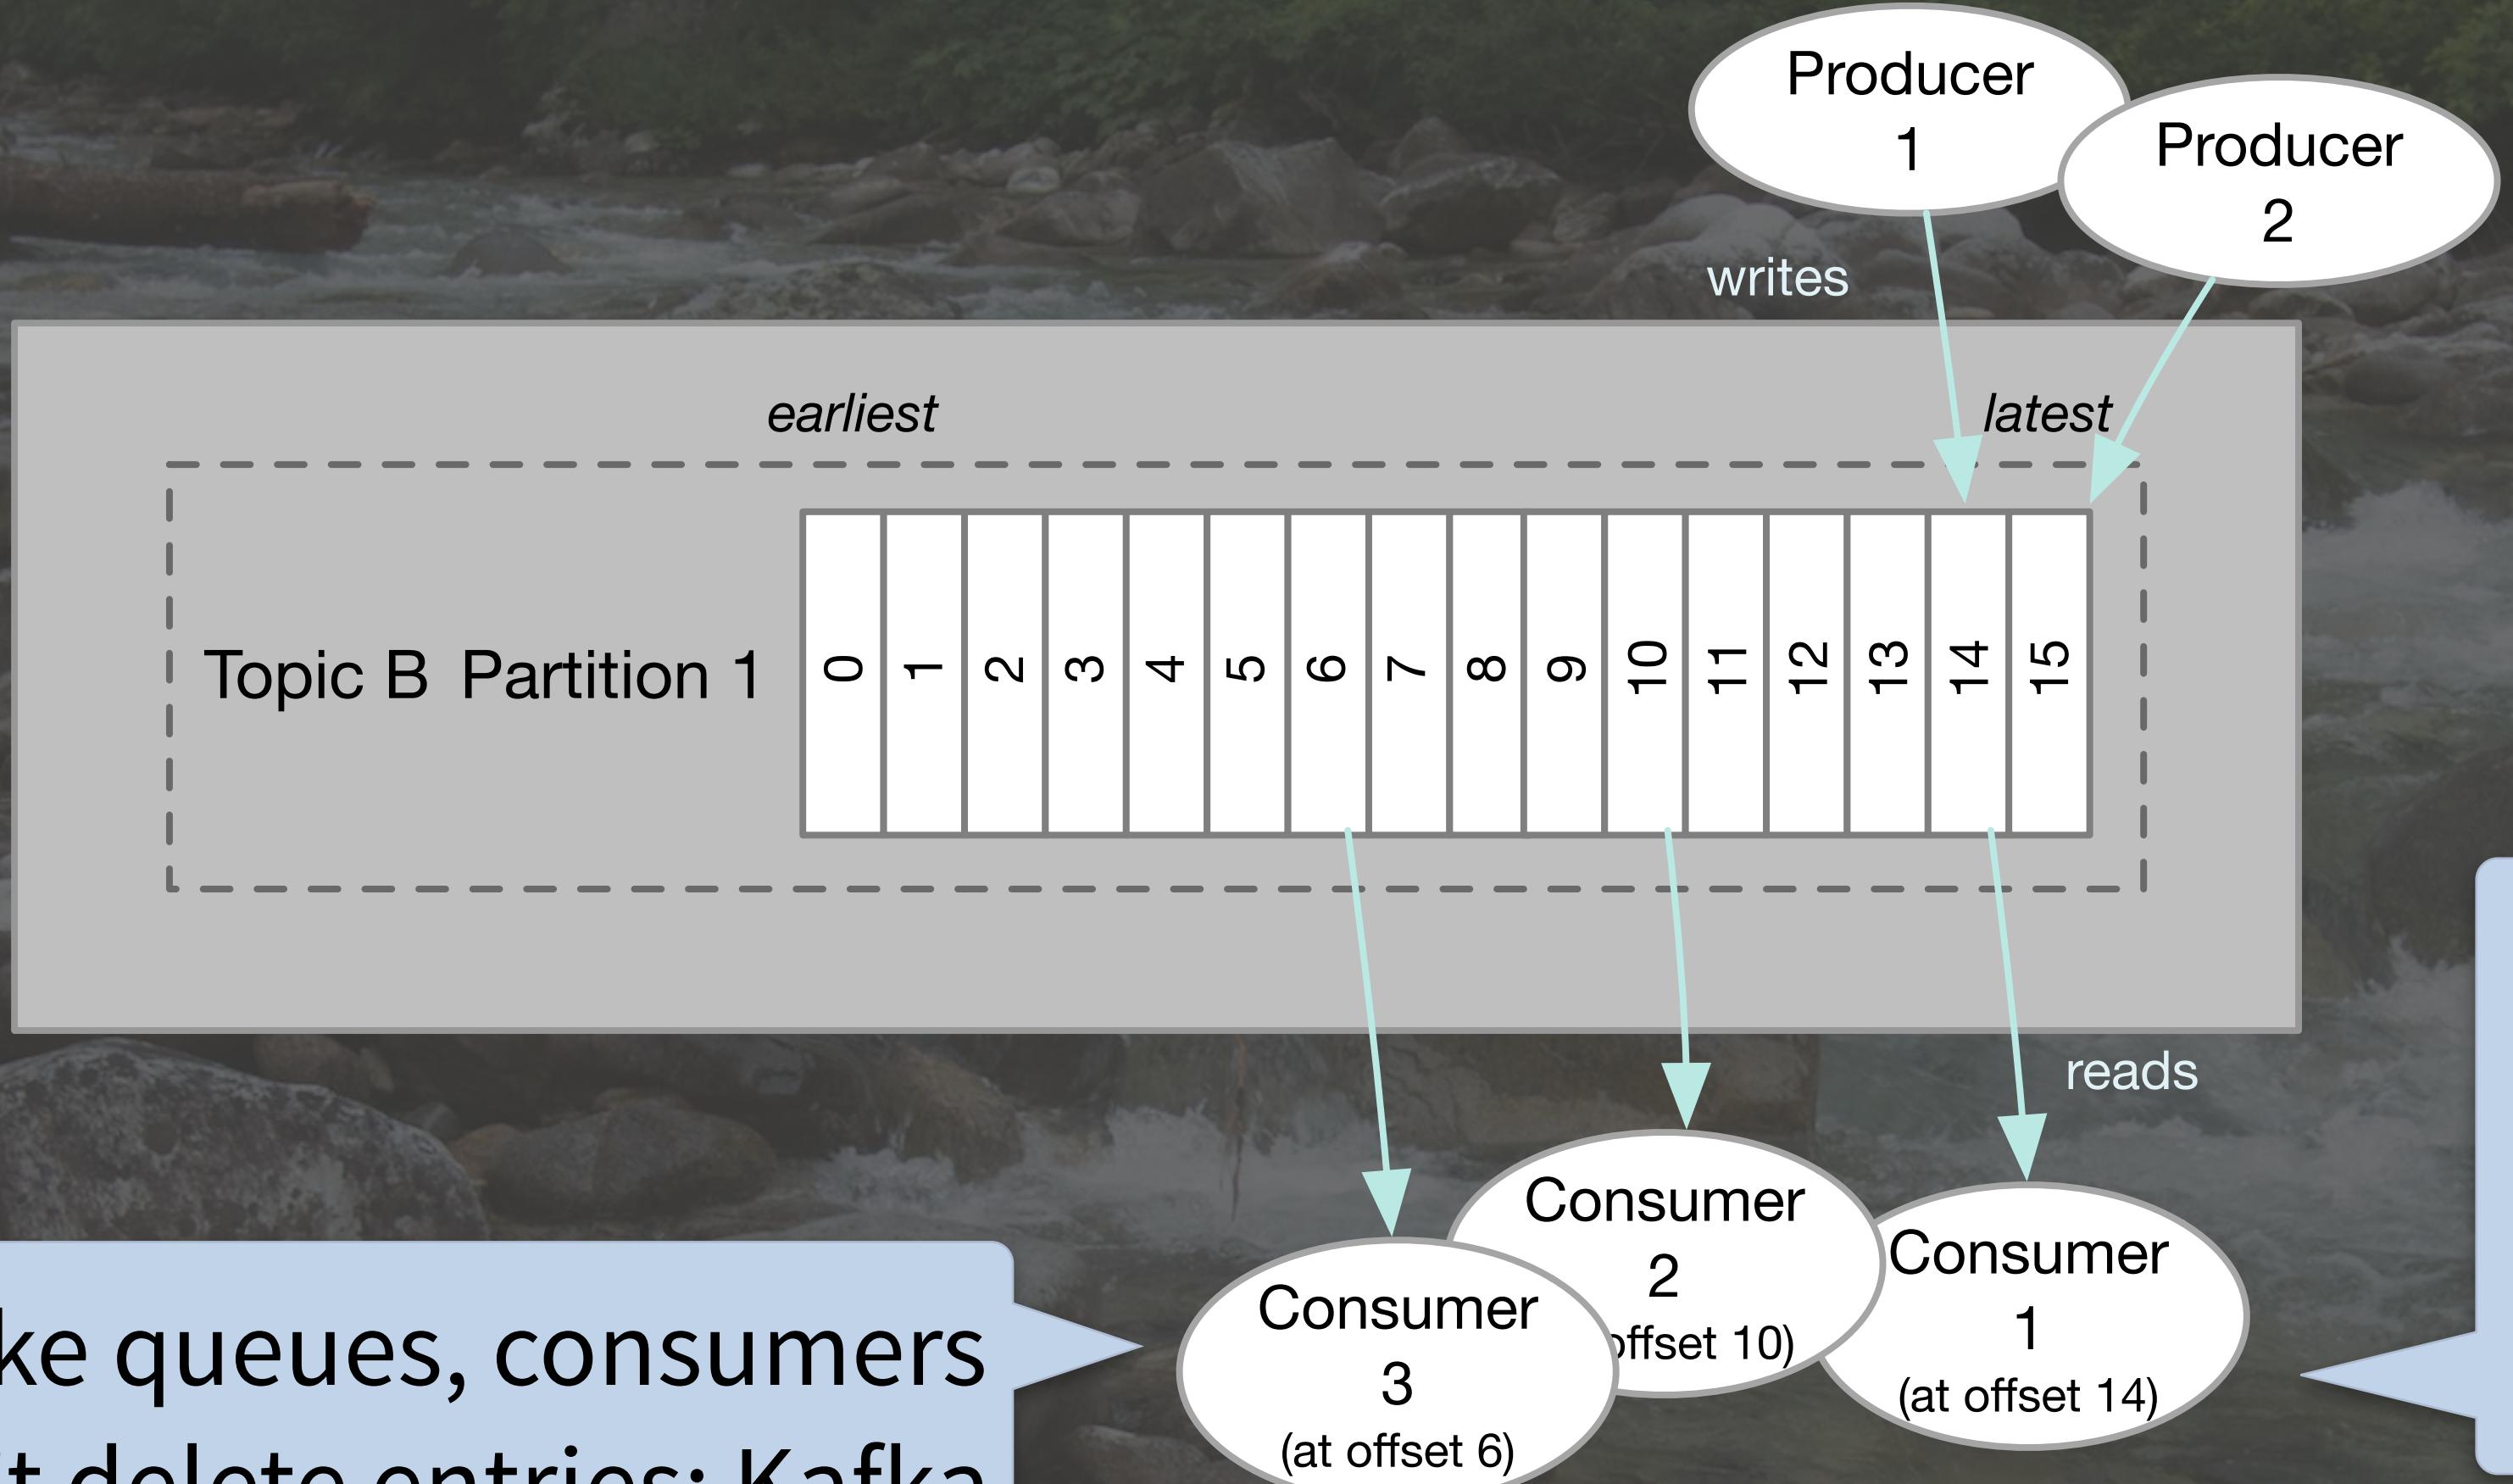
Organized into topics

Topics are partitioned, replicated, and distributed



Why Kafka?

Logs, not queues!

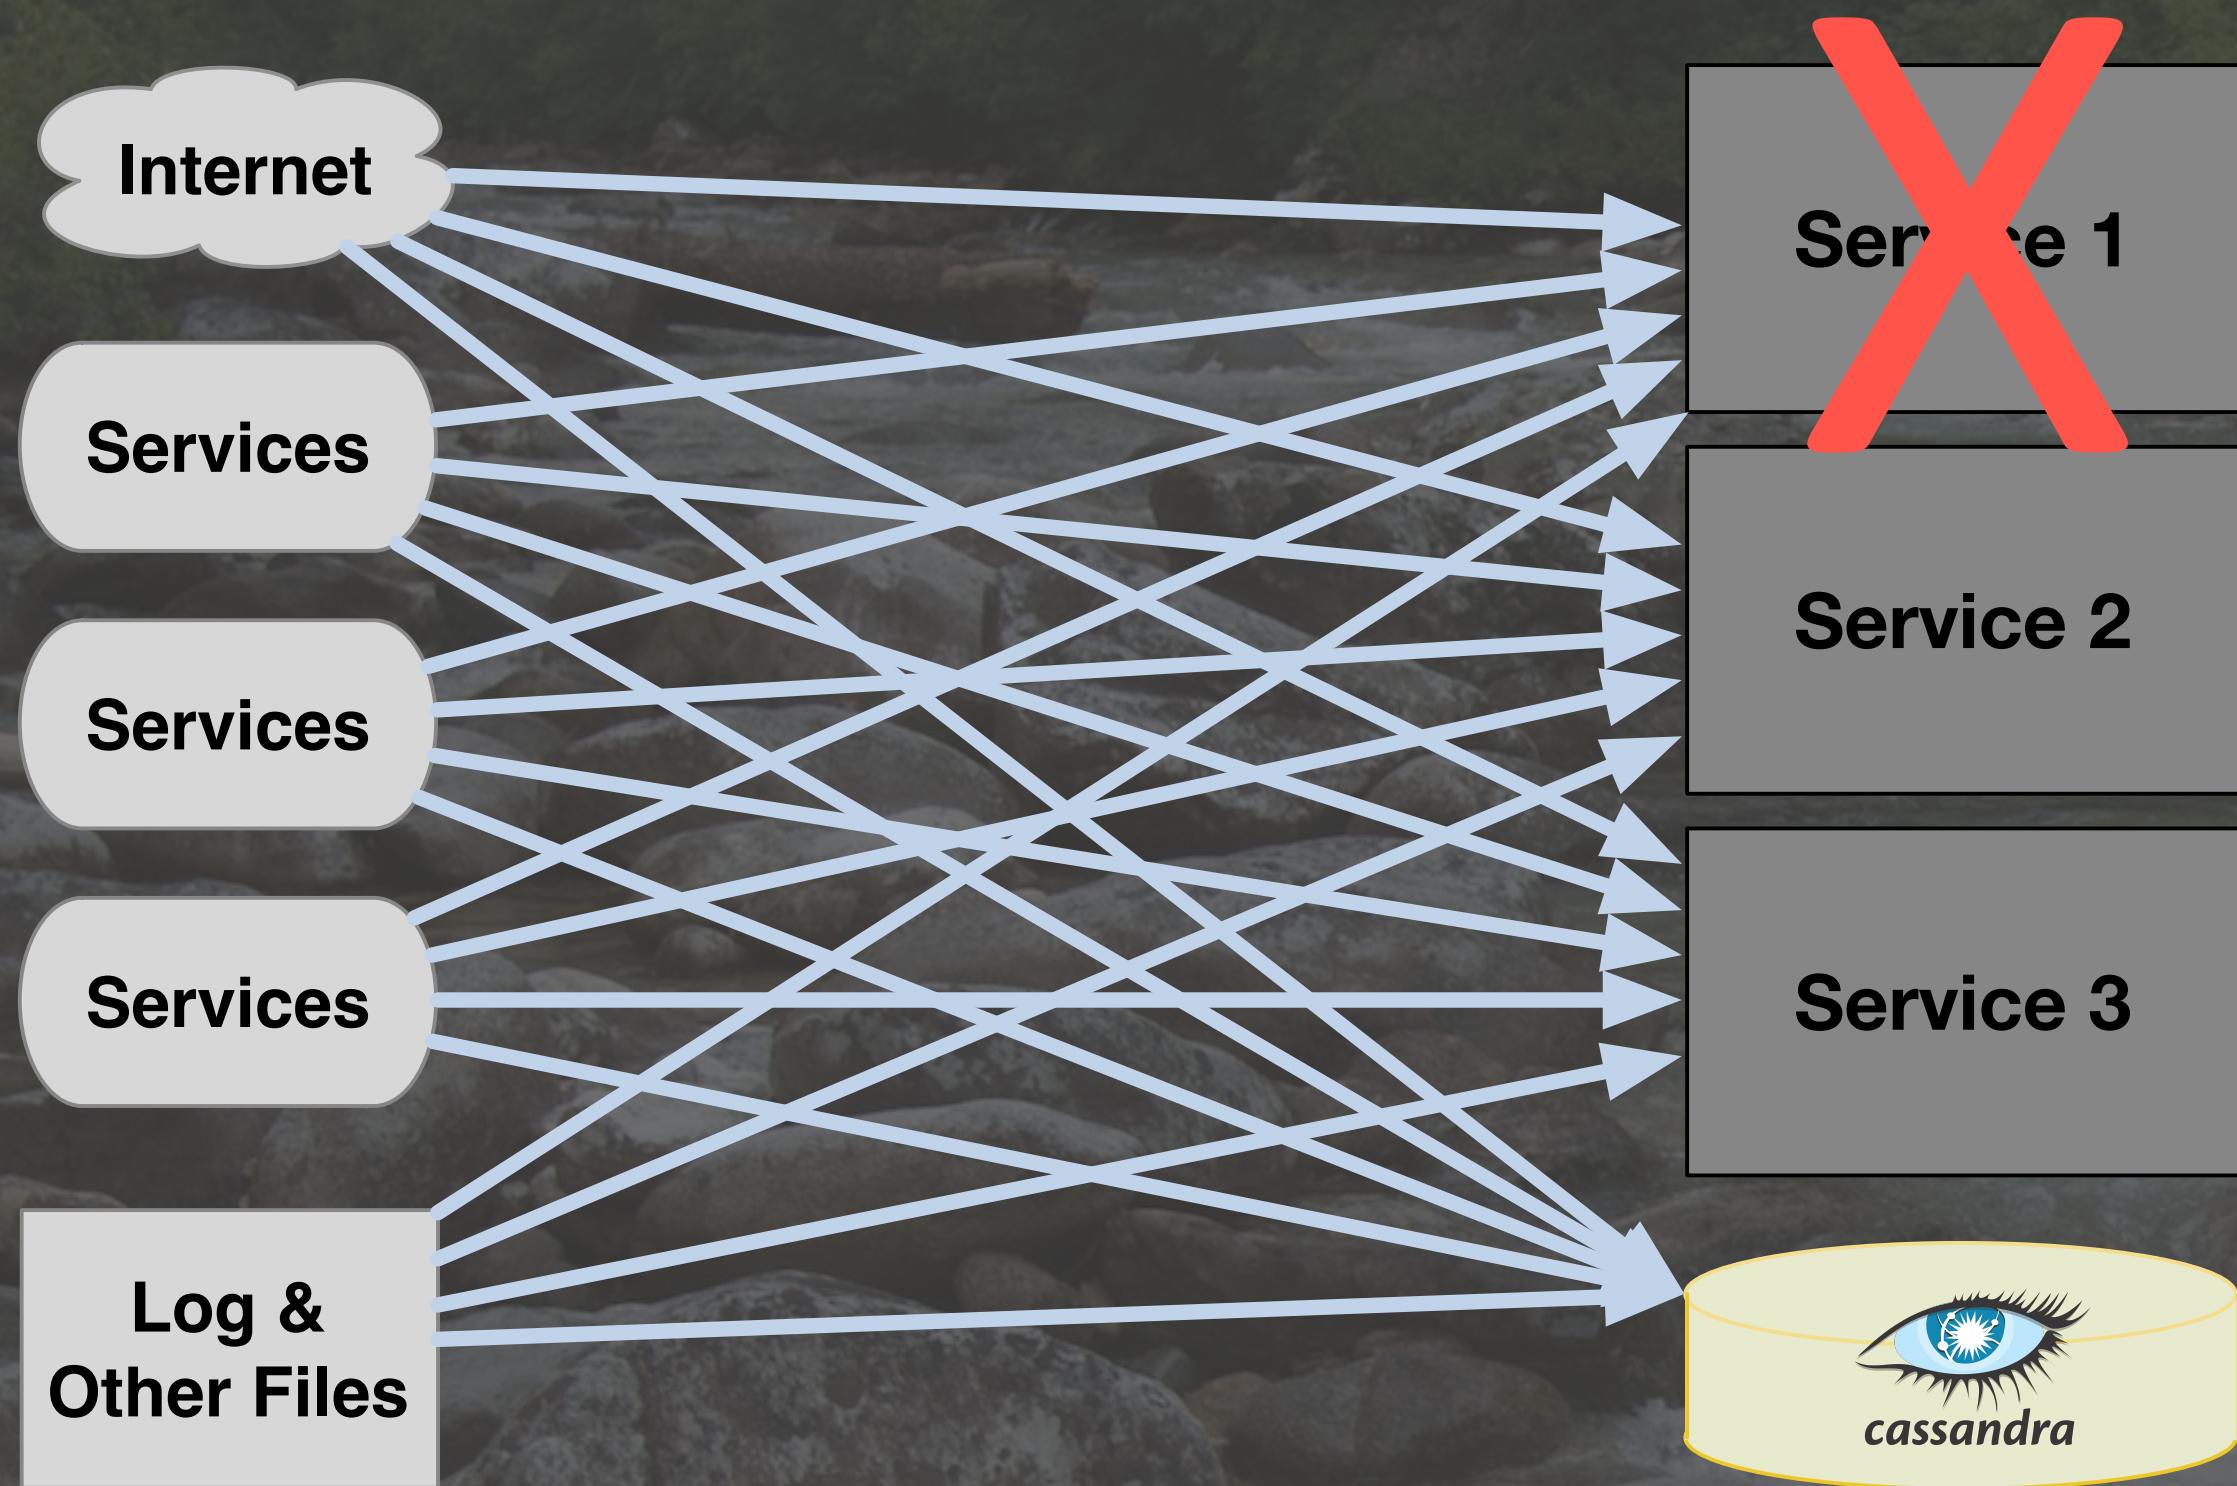


Unlike queues, consumers don't delete entries; Kafka manages their lifecycles



Using Kafka

Before:



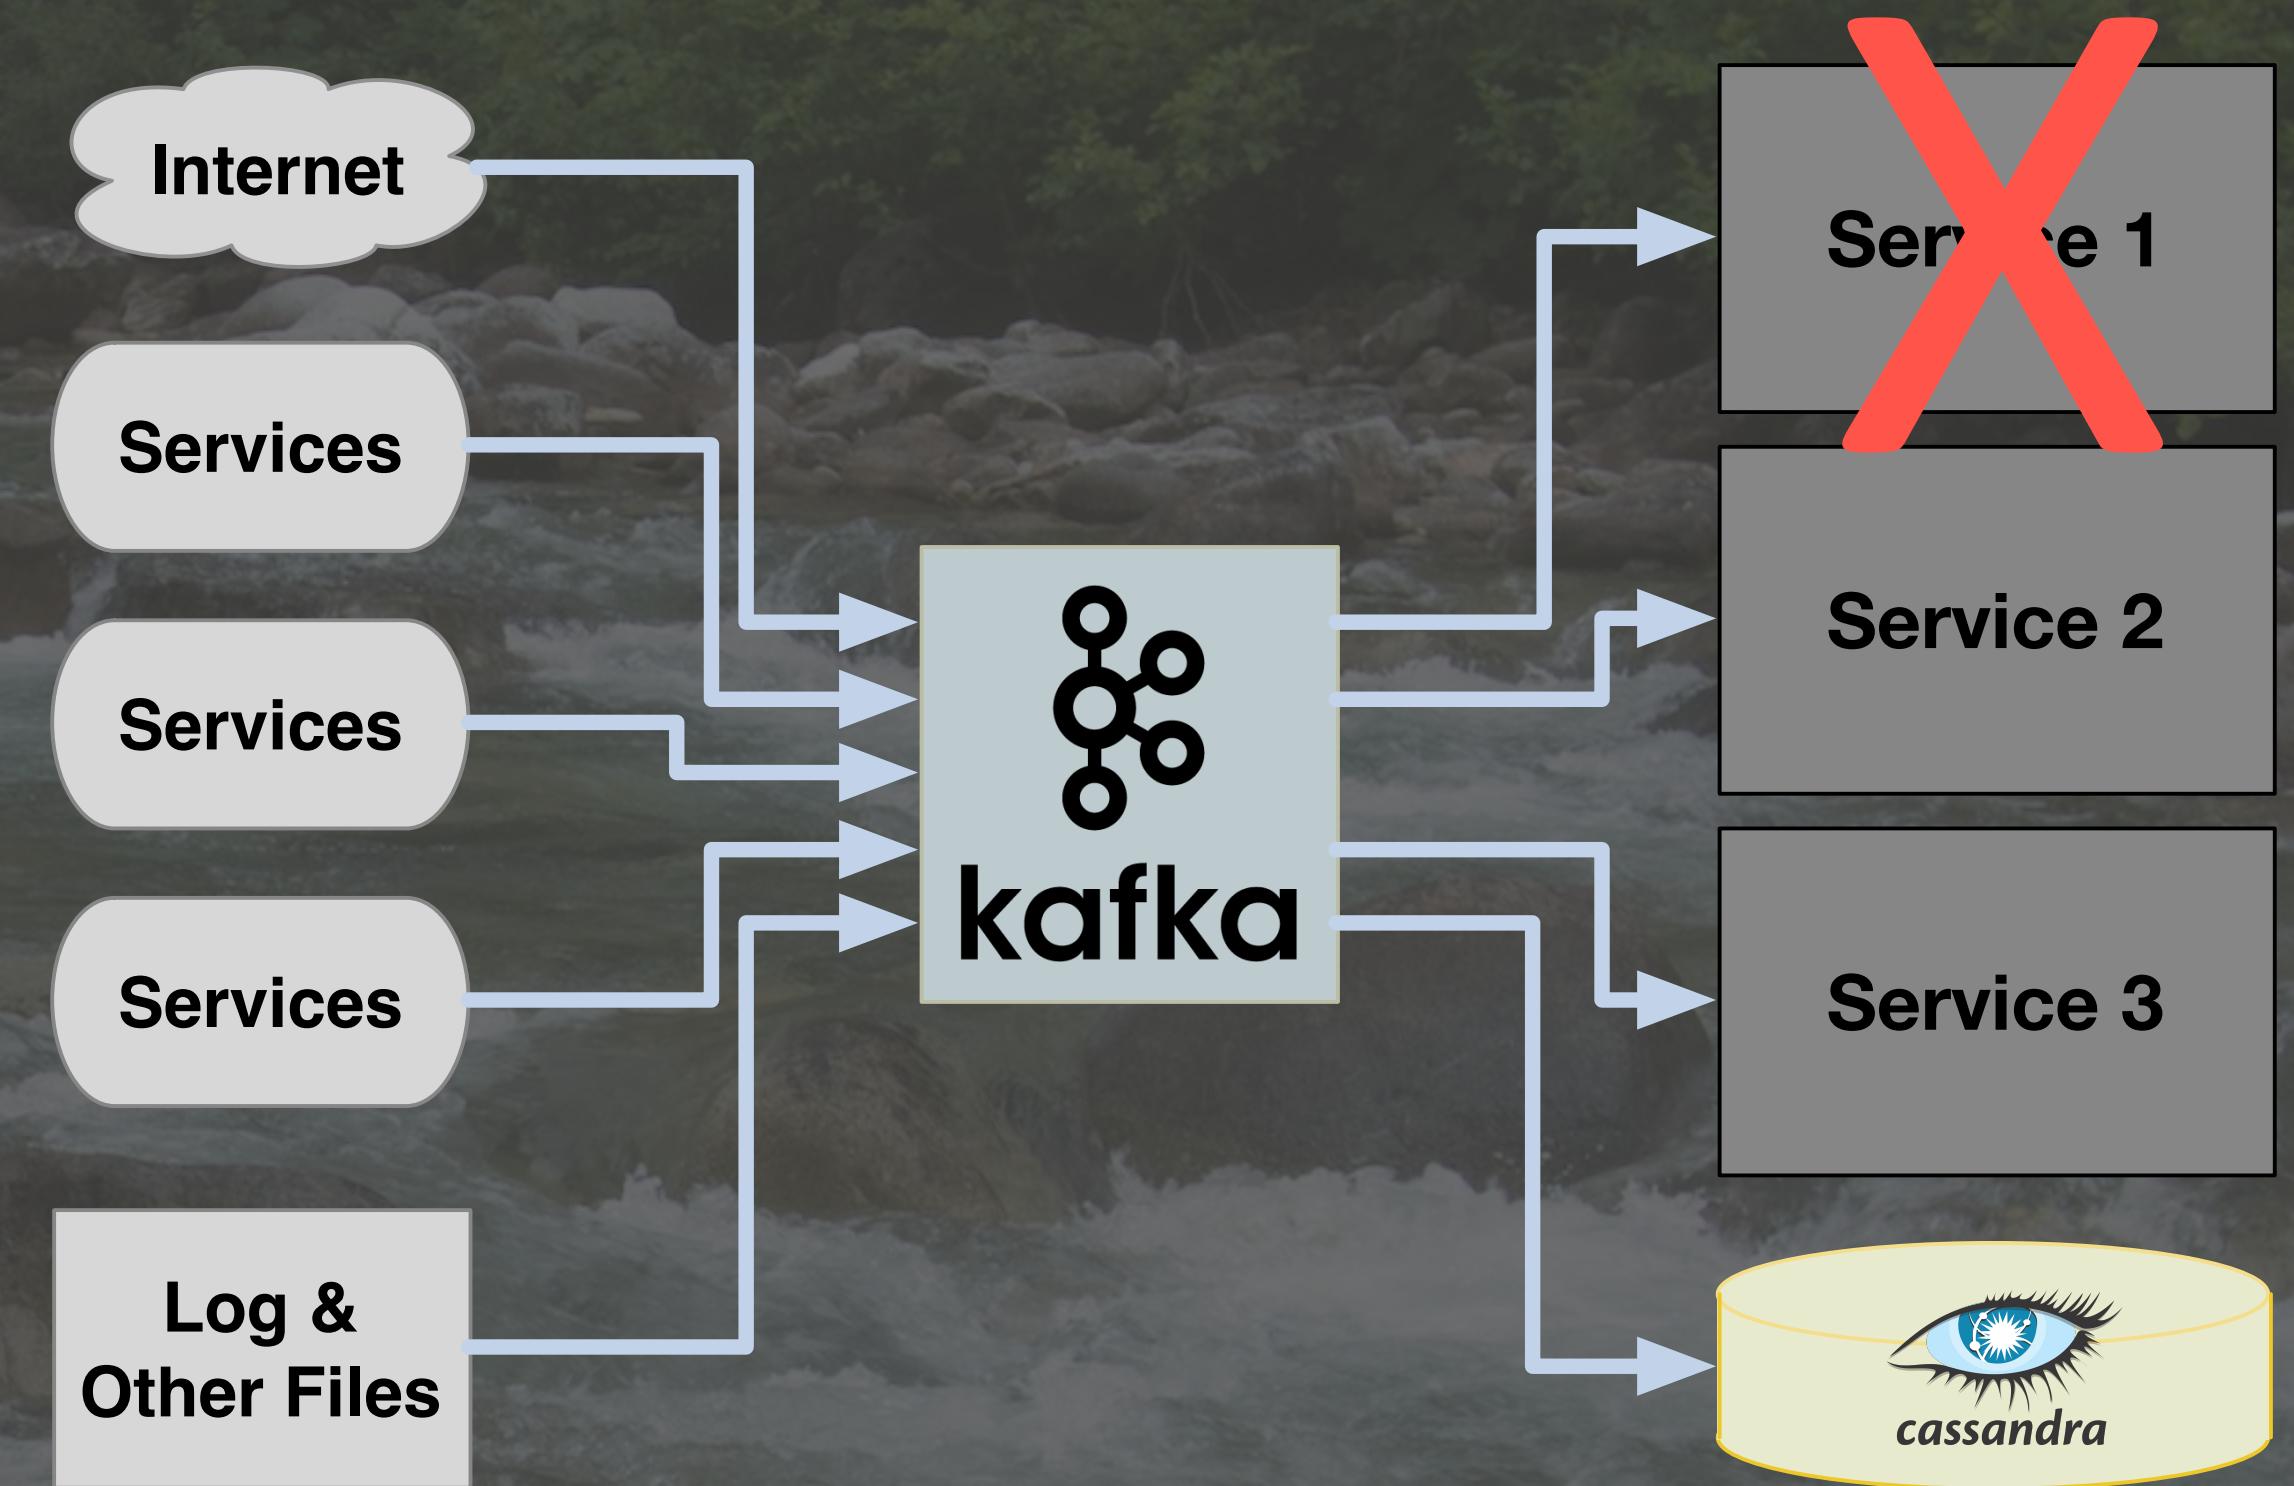
Producers

$N * M$ links

Consumers

 Messy and fragile;
what if “Service 1”
goes down?

After:

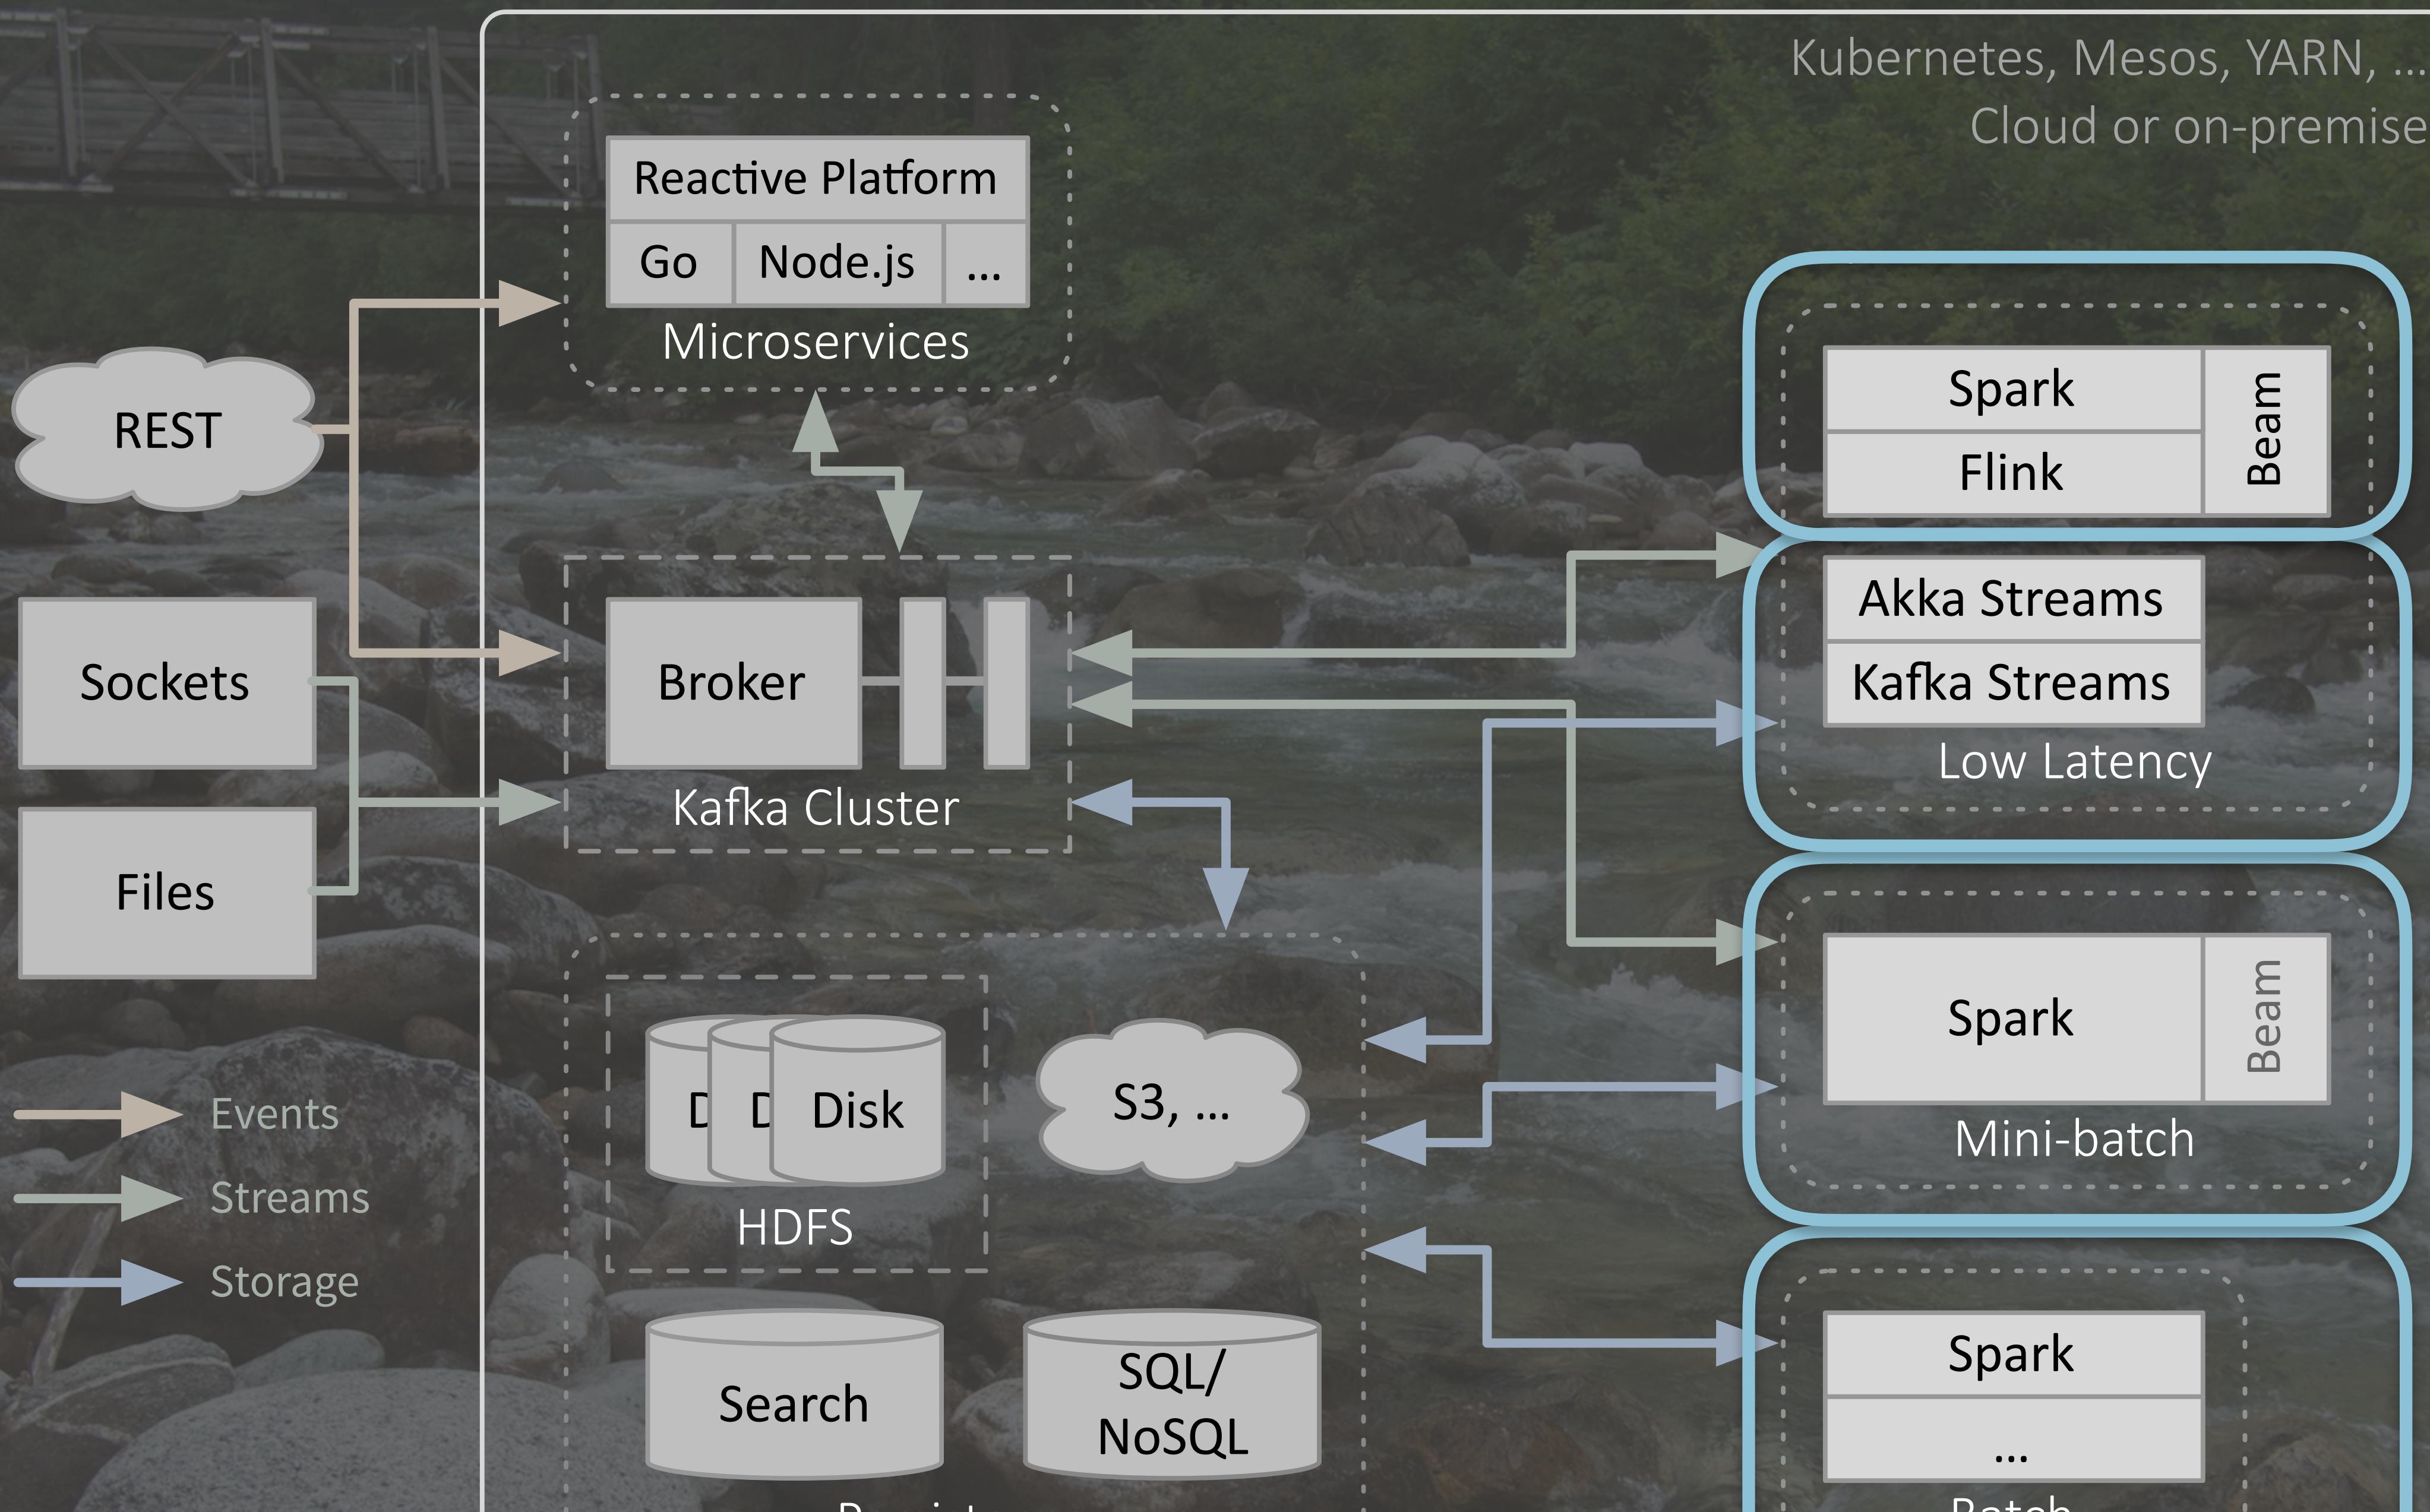


Producers

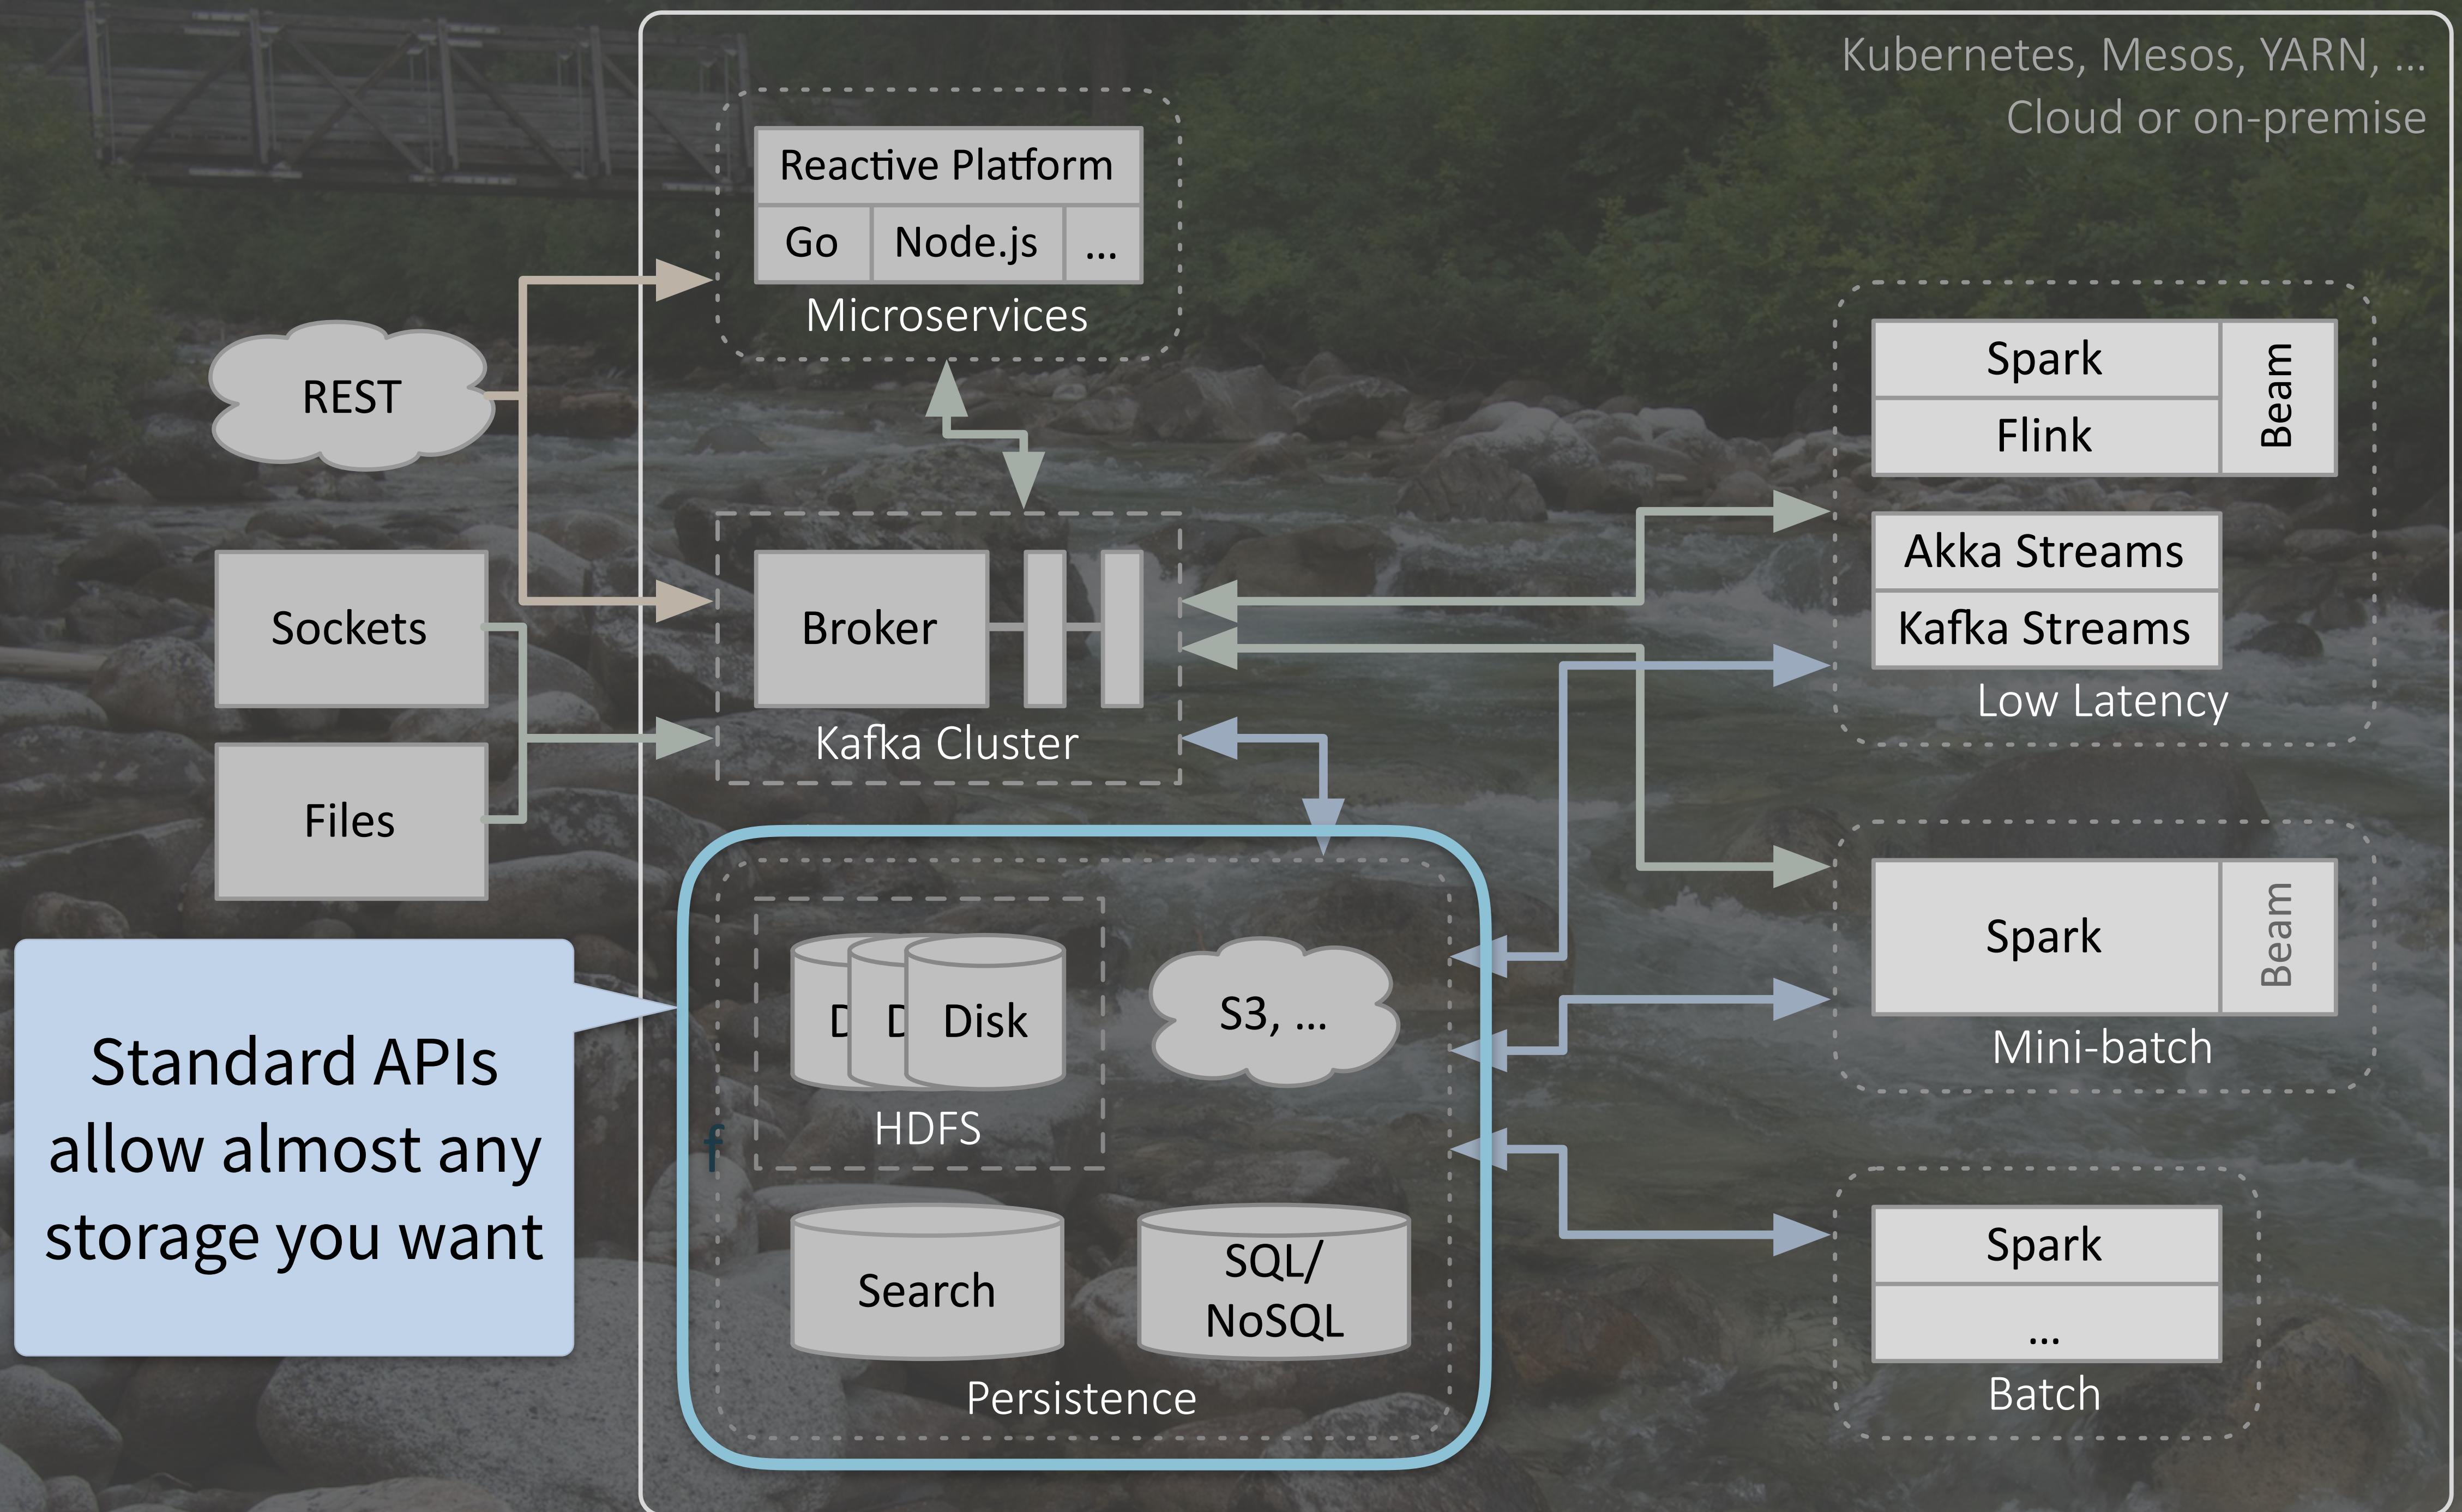
$N + M$ links

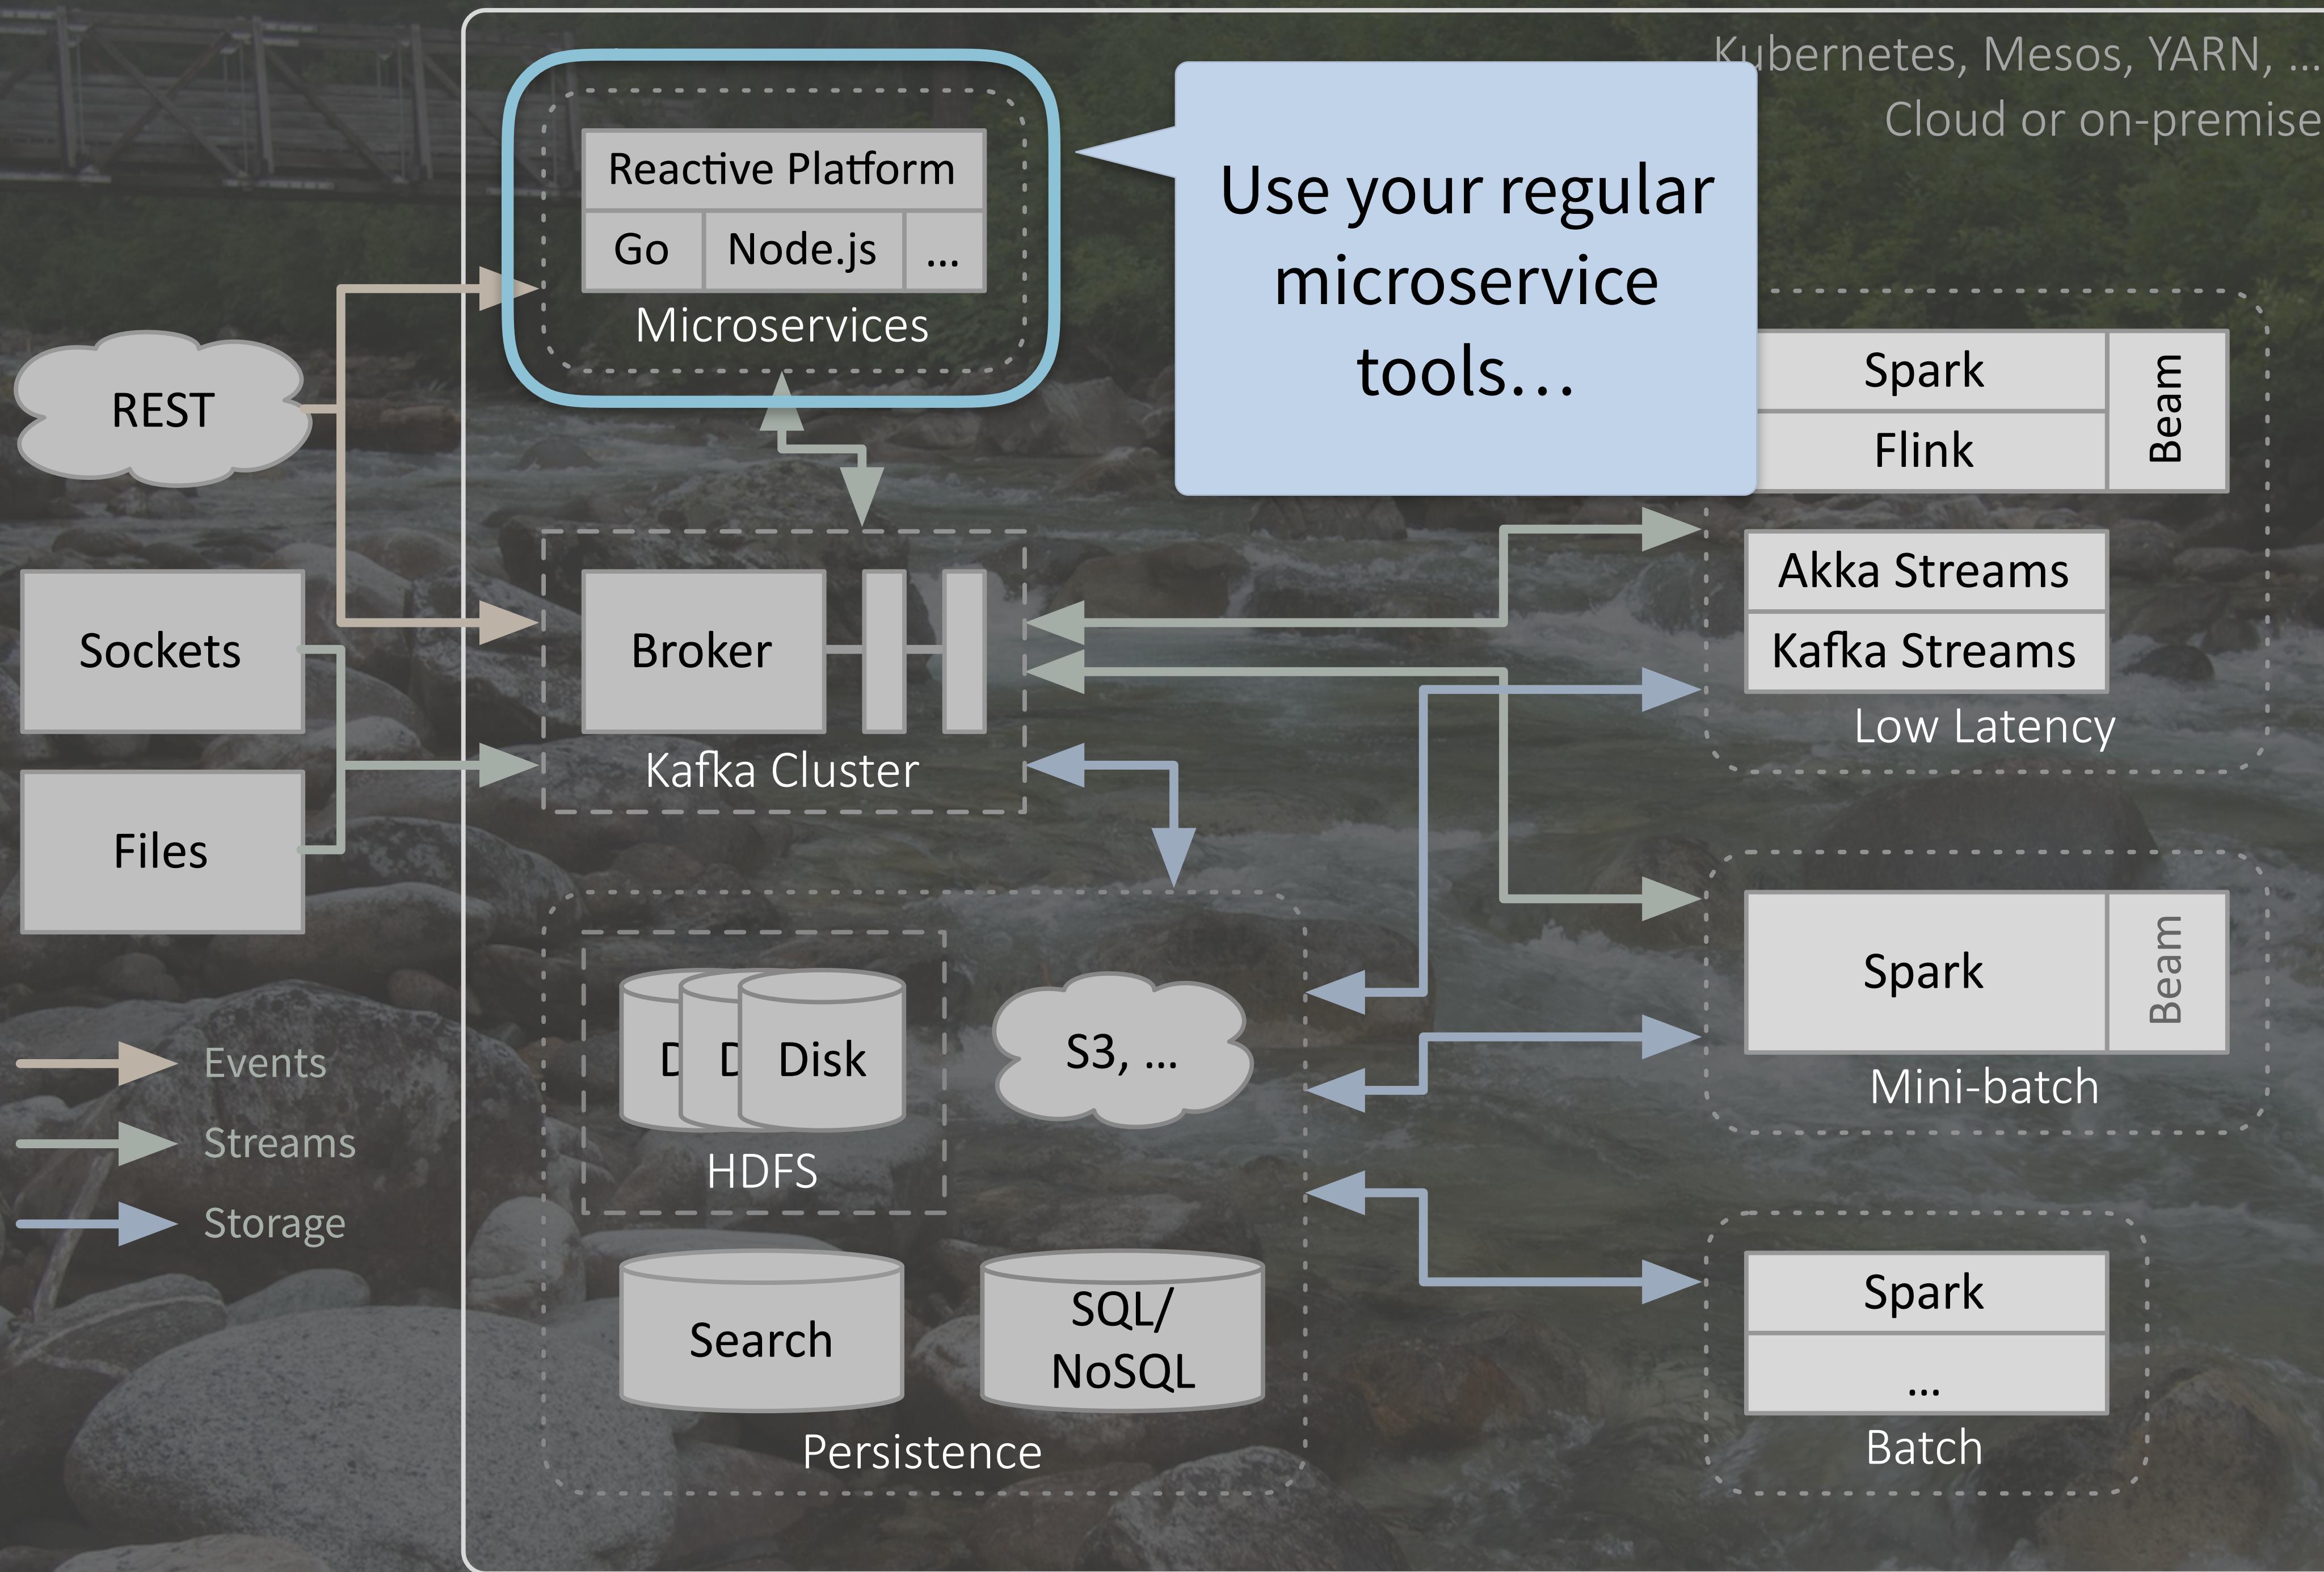
Consumers

Simpler and more
robust! Loss of Service
1 means no data loss.



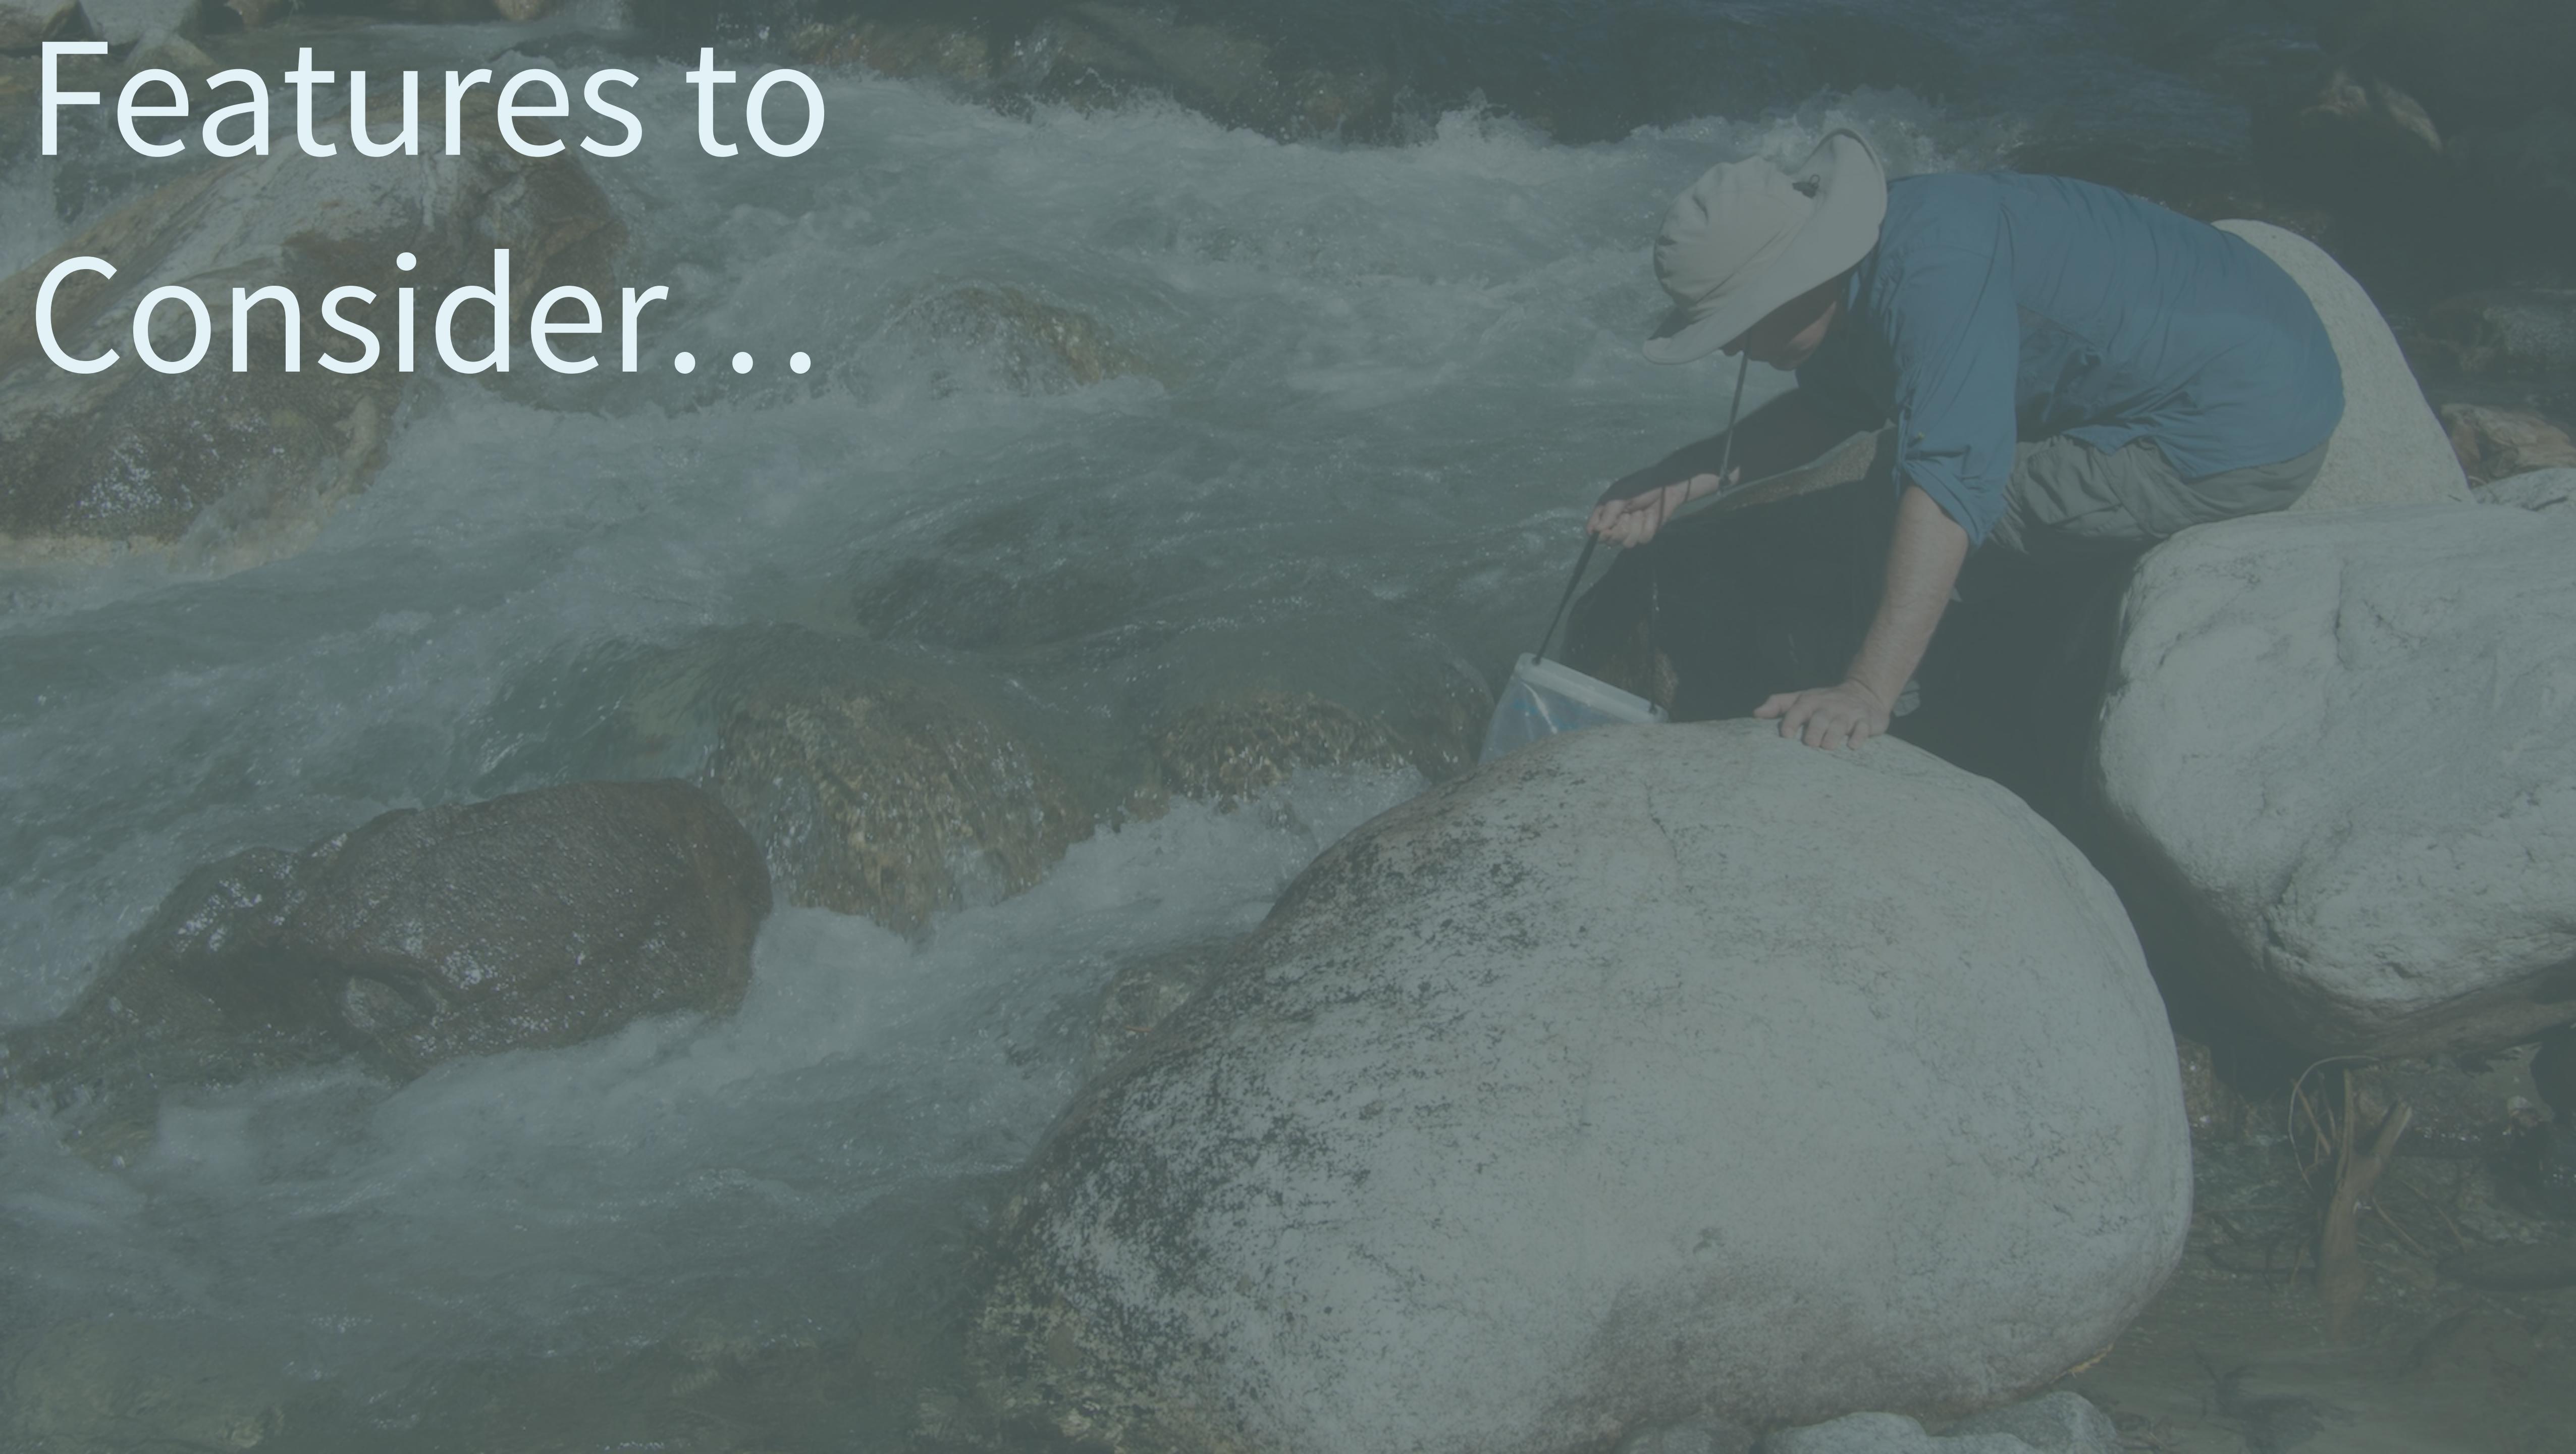
Lots of streaming engine options... too many.







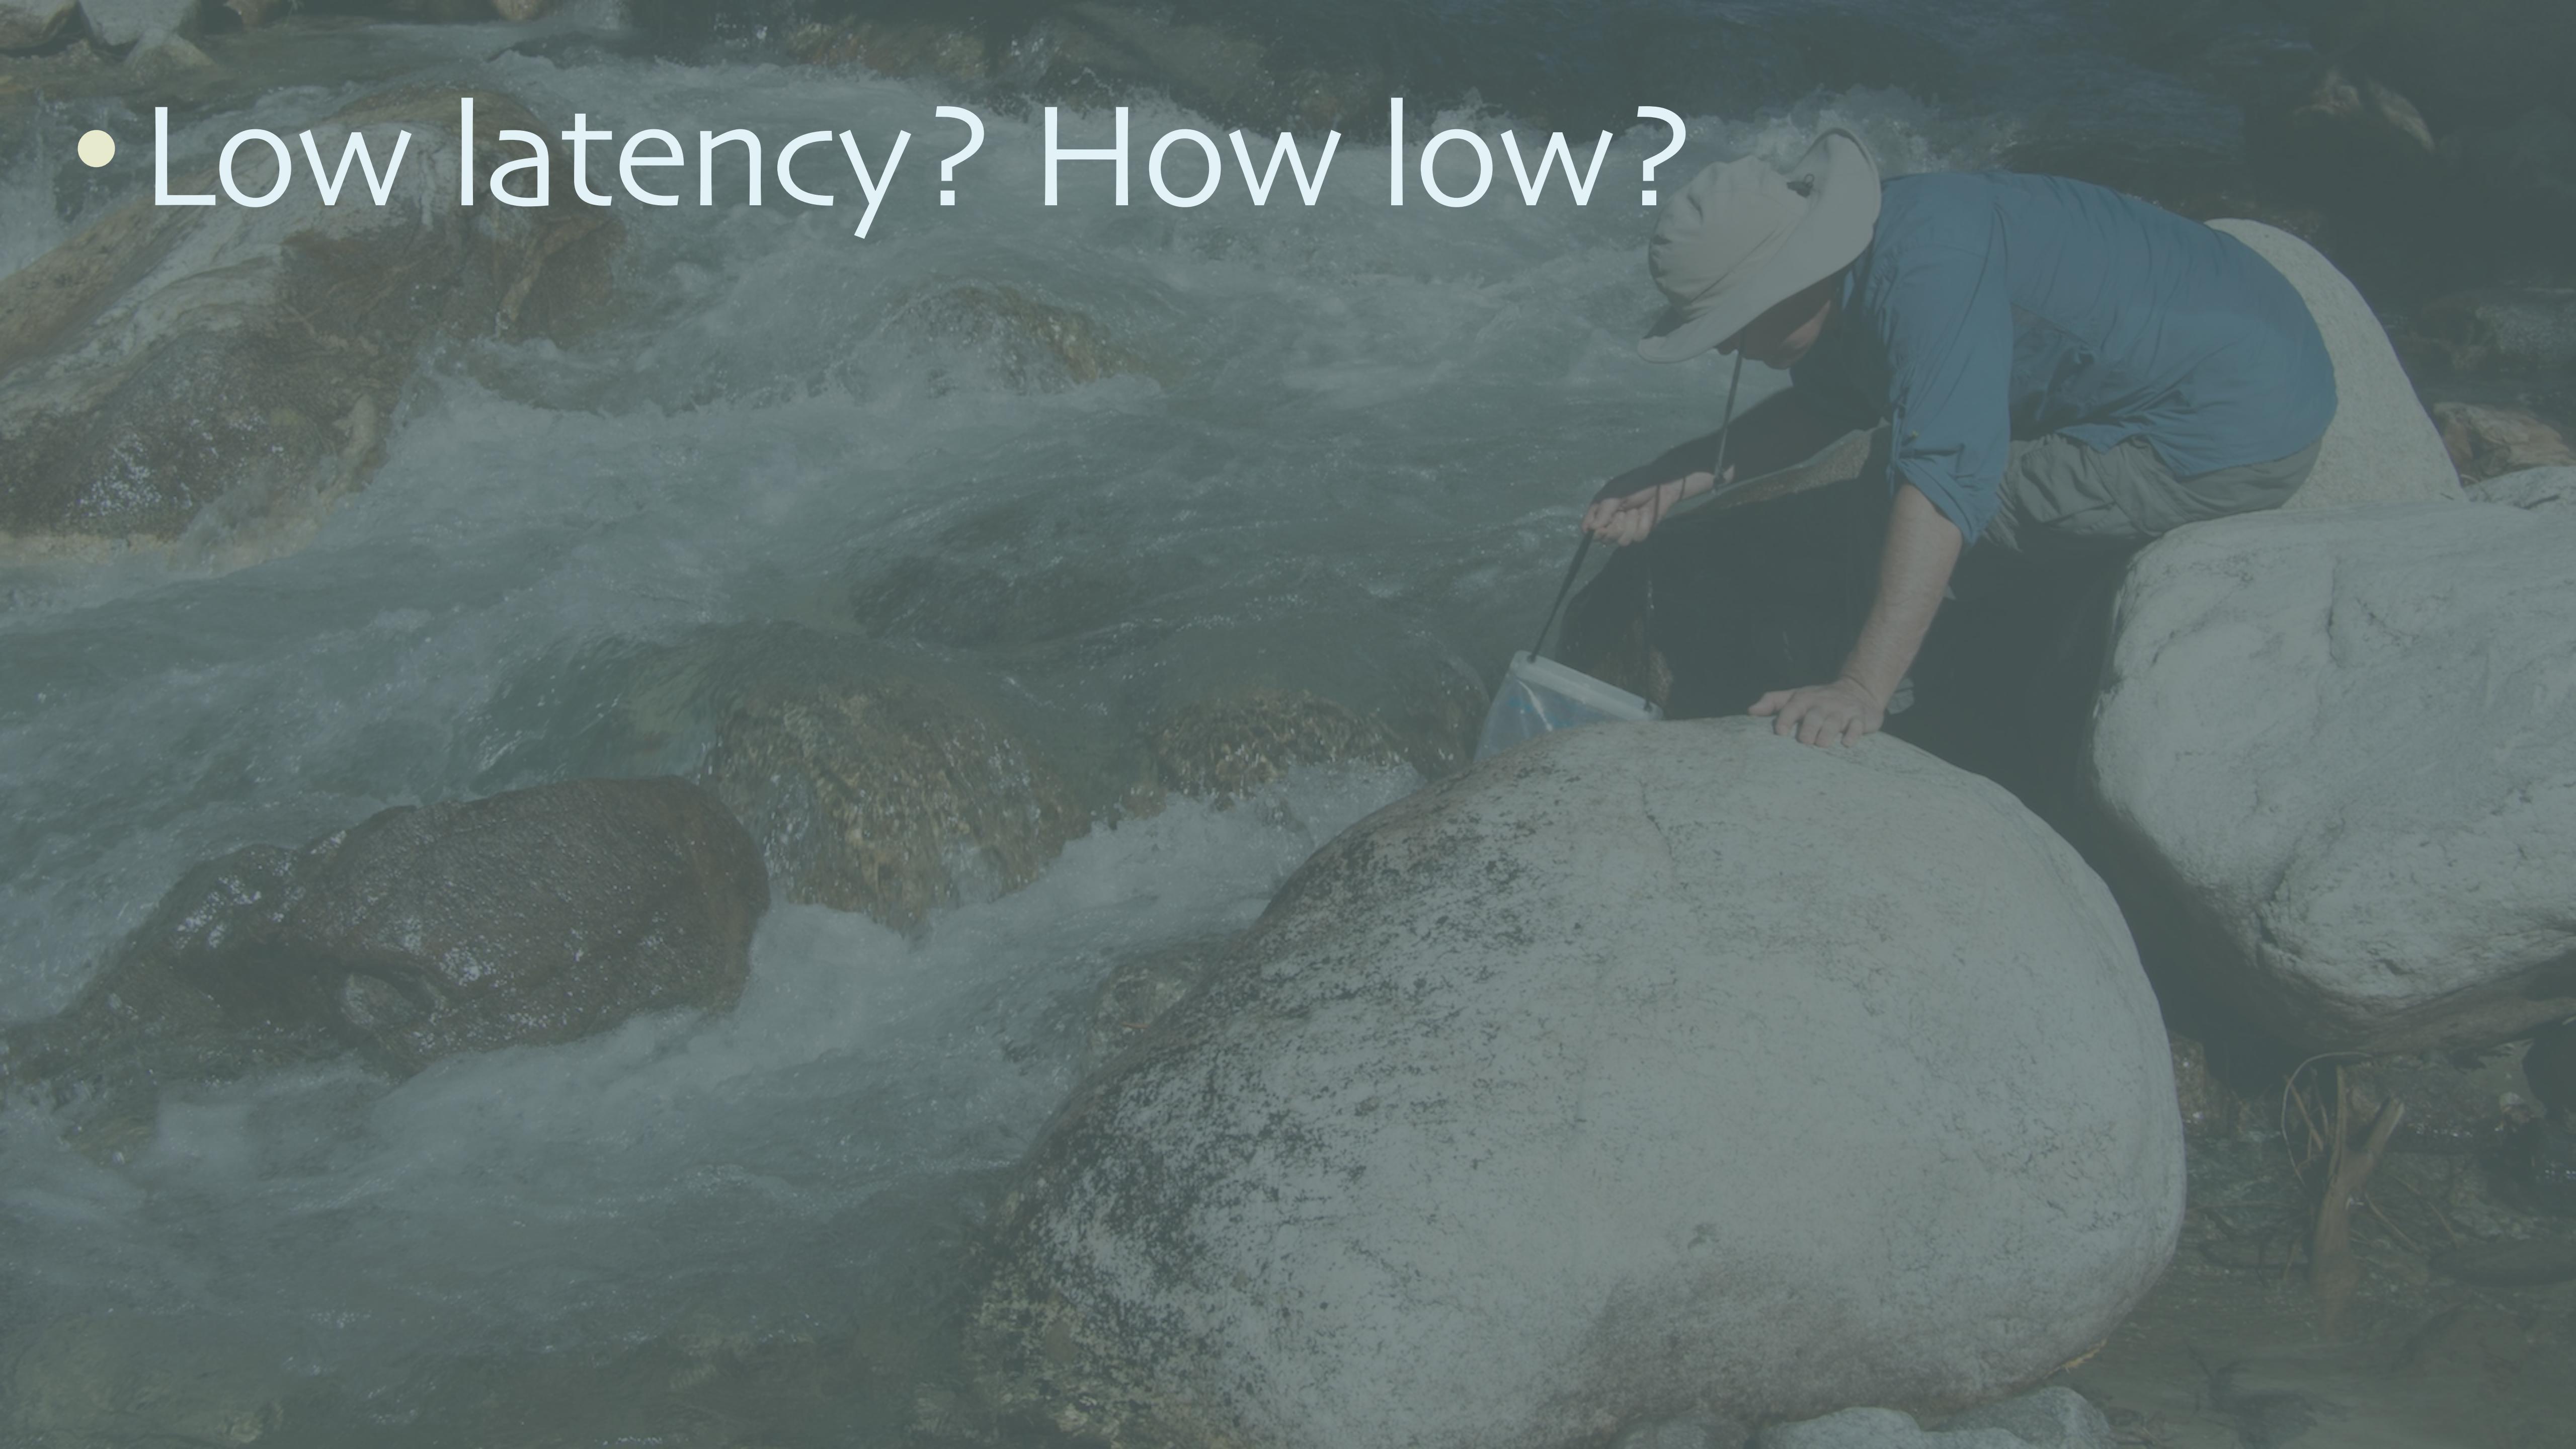
Streaming Engines



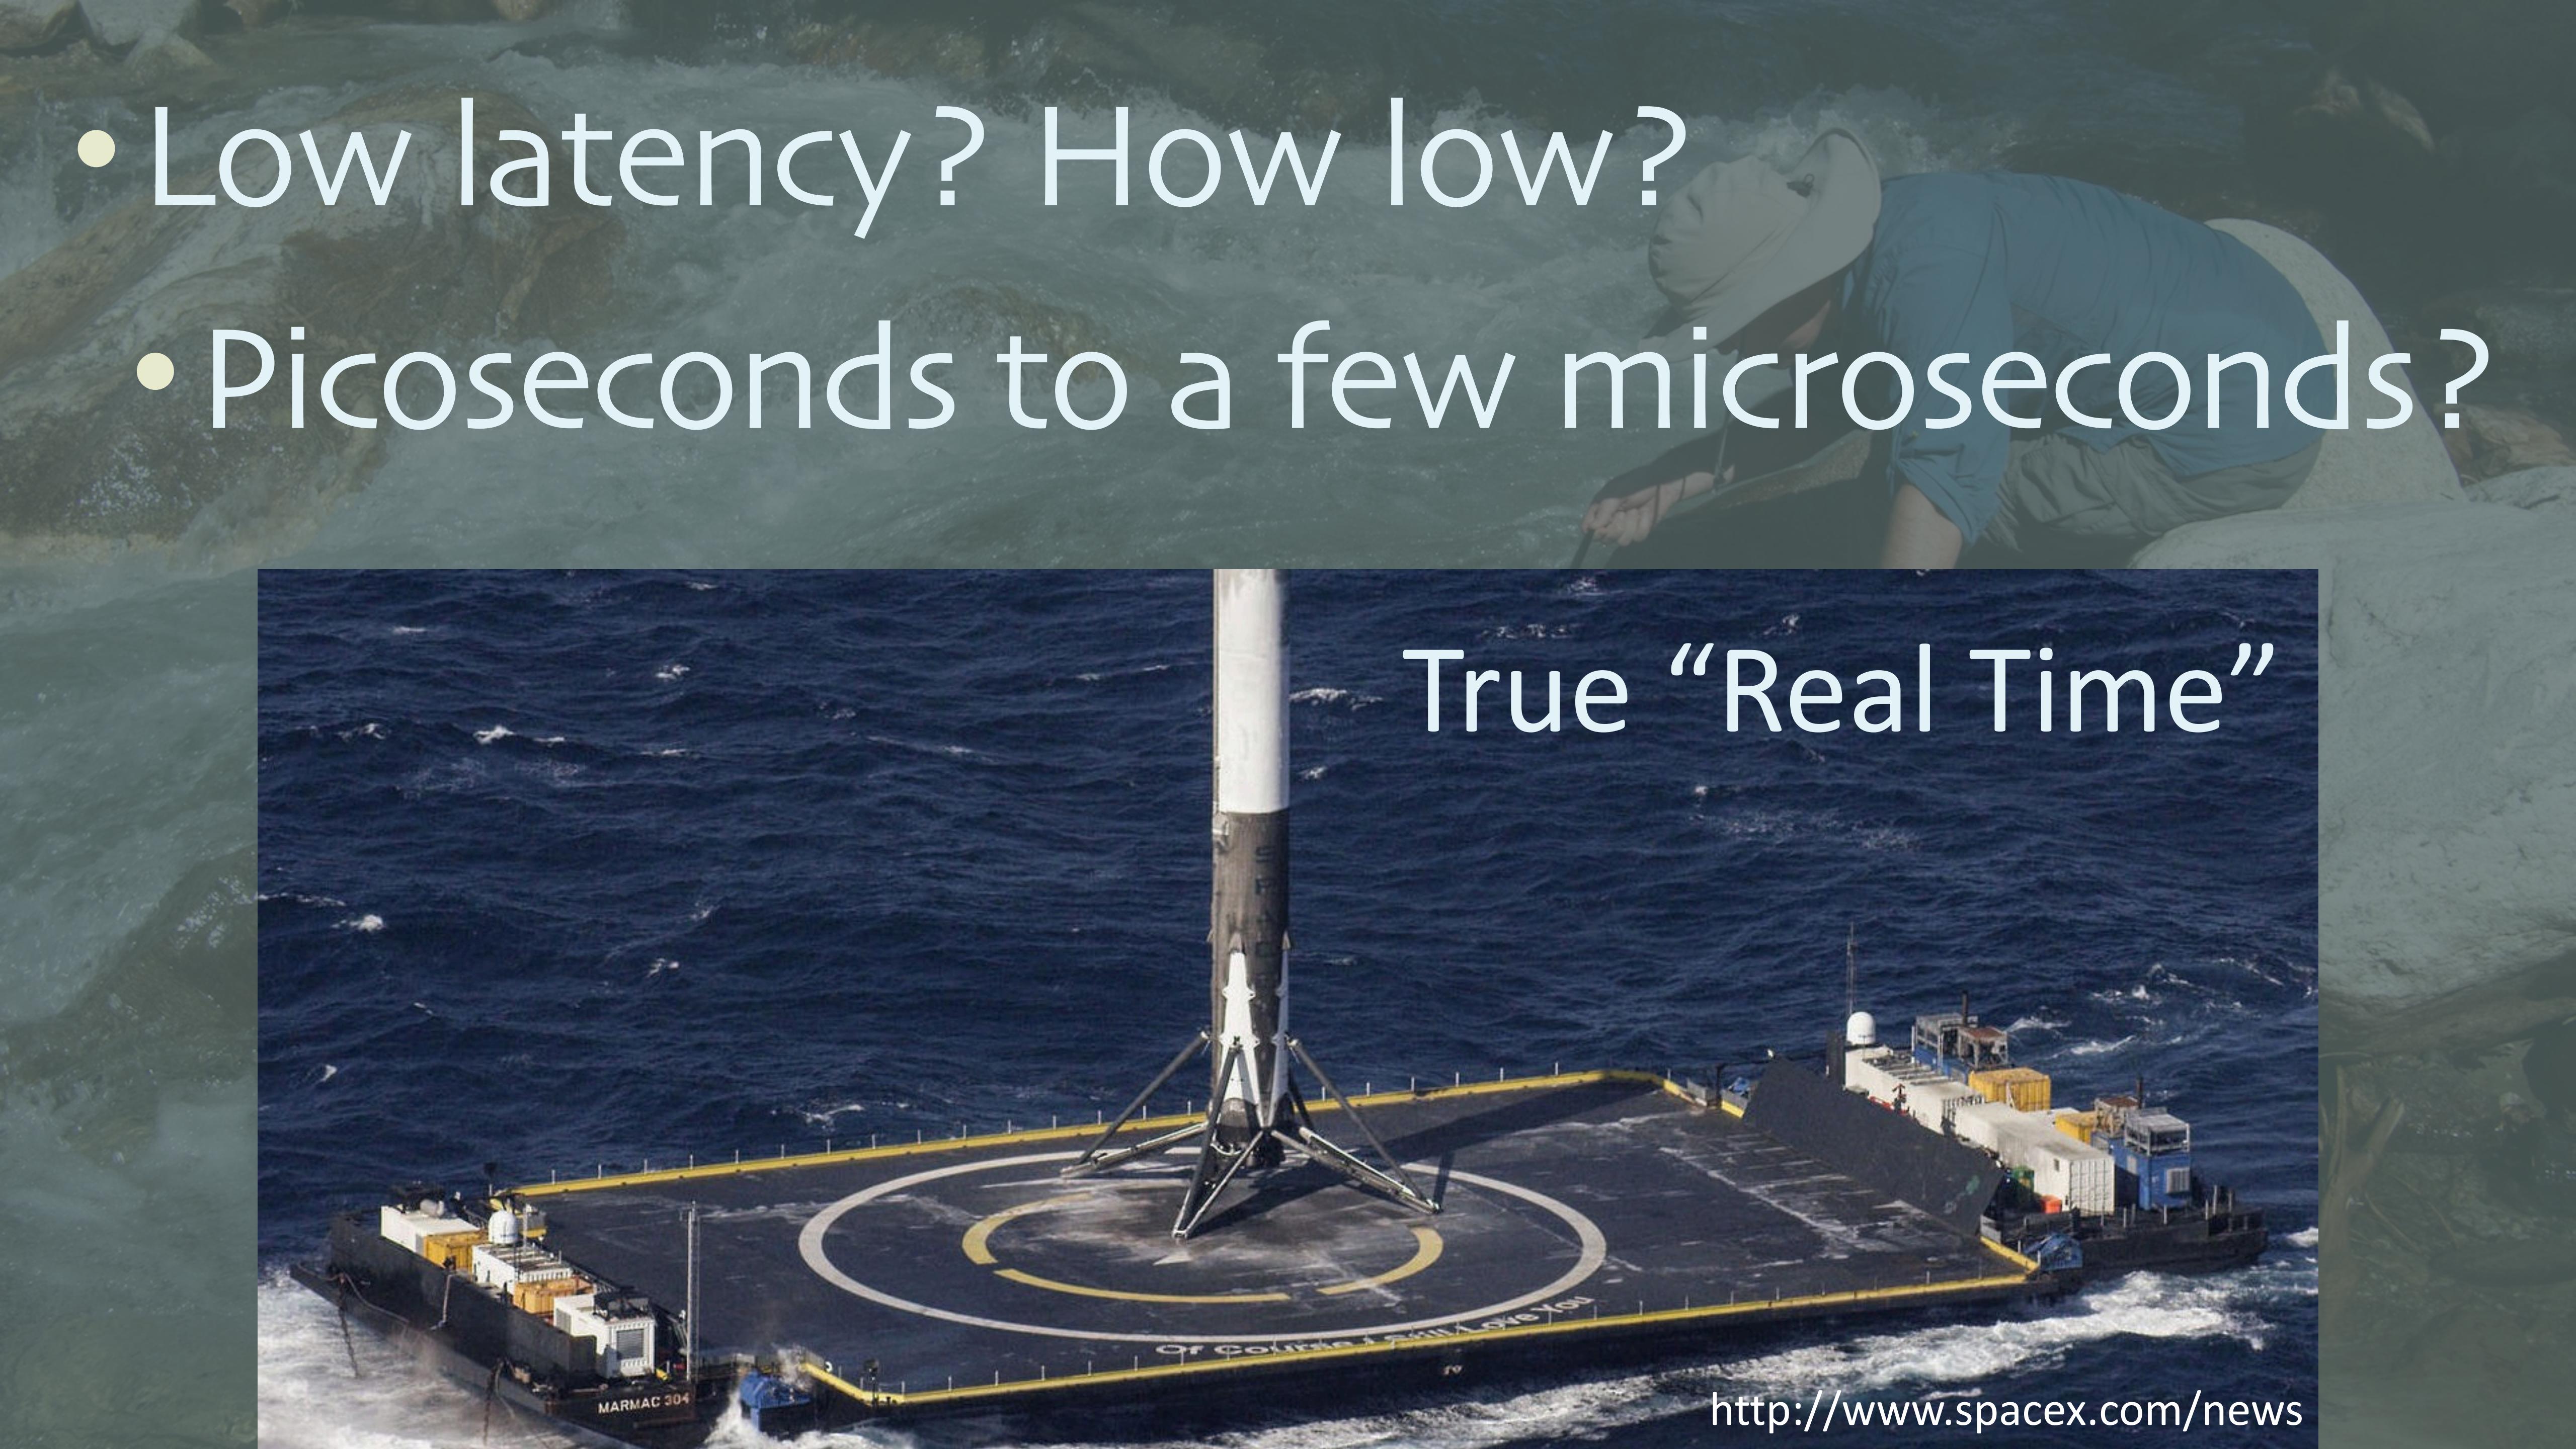
Features to Consider...

- Low latency? How low?
- High Volume: How high?
- Which kinds of data processing?
- Process data individually or in bulk?
- Preferred application architecture and DevOps processes?
- Integration with other services

- Low latency? How low?



- Low latency? How low?
- Picoseconds to a few microseconds?



True “Real Time”



- Low latency? How low?
 - Picoseconds to a few microseconds?
 - Custom hardware (FPGAs).
 - “Kernel bypass” network HW/SW.
 - Custom C++ code.

- Low latency? How low?
- < 100 microseconds?



- Low latency? How low?
- < 100 microseconds?
 - Fast JVM message handlers.
 - Akka Actors
 - LMAX Disruptor

- Low latency? How low?
- < 10 milliseconds?

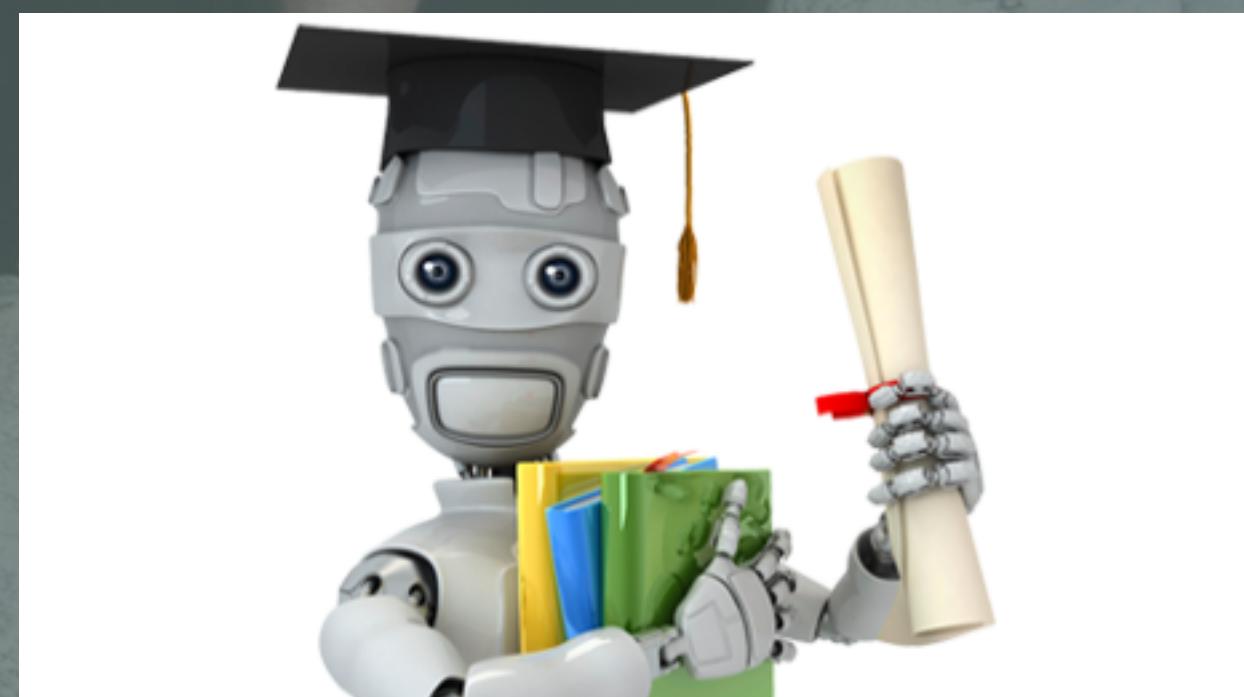


- Low latency? How low?
- < 10 milliseconds?
- Fast data streaming tools like Flink and more recently Spark, Akka (and Akka Streams), and Kafka Streams.

- Low latency? How low?
- < hundreds of milliseconds?



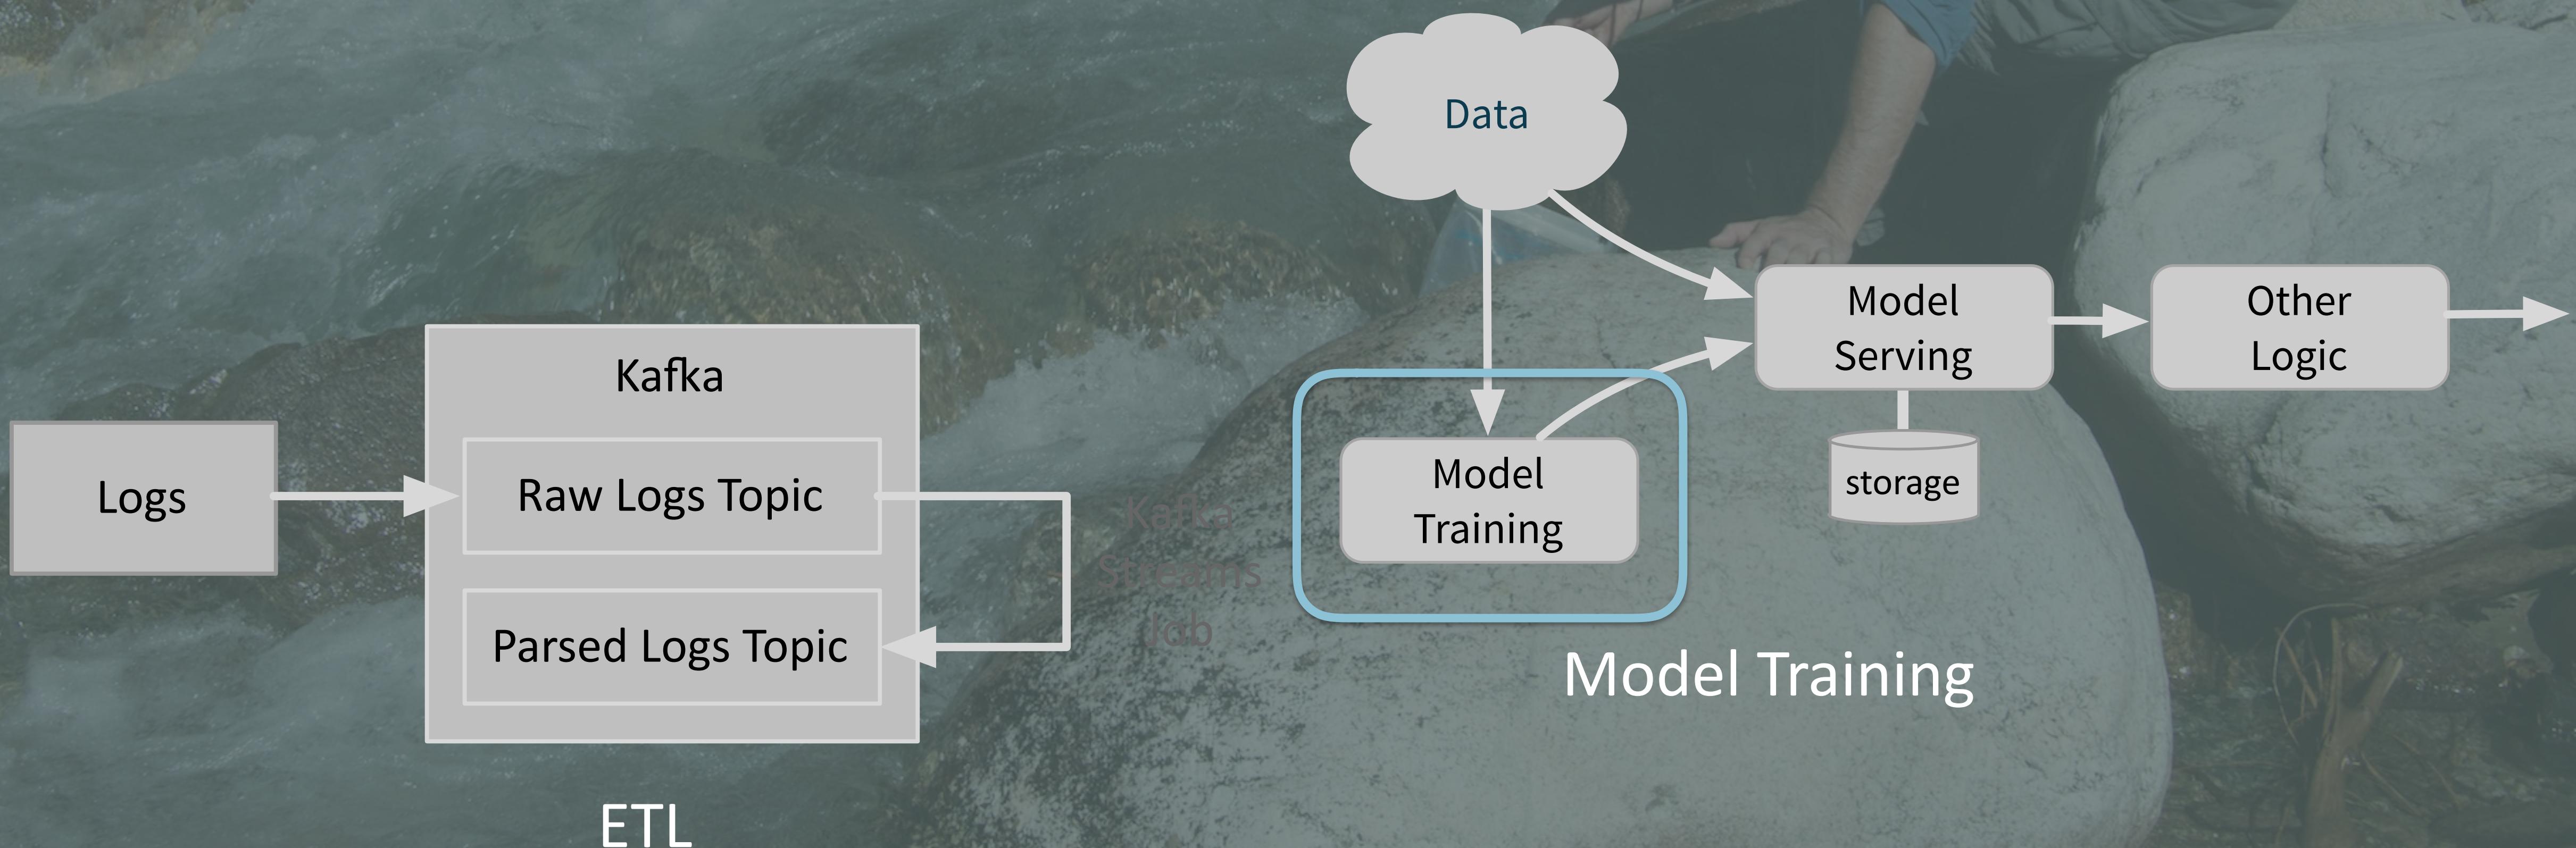
<https://github.com/keen/dashboards>



<https://www.coursera.org/learn/machine-learning>

- Low latency? How low?
- < hundreds of milliseconds?
- “micro-batches”
- Processing records in bulk, e.g.,
Spark’s micro-batch model and
“streaming SQL” over windows.

- Low latency? How low?
- < 1 second to minutes?

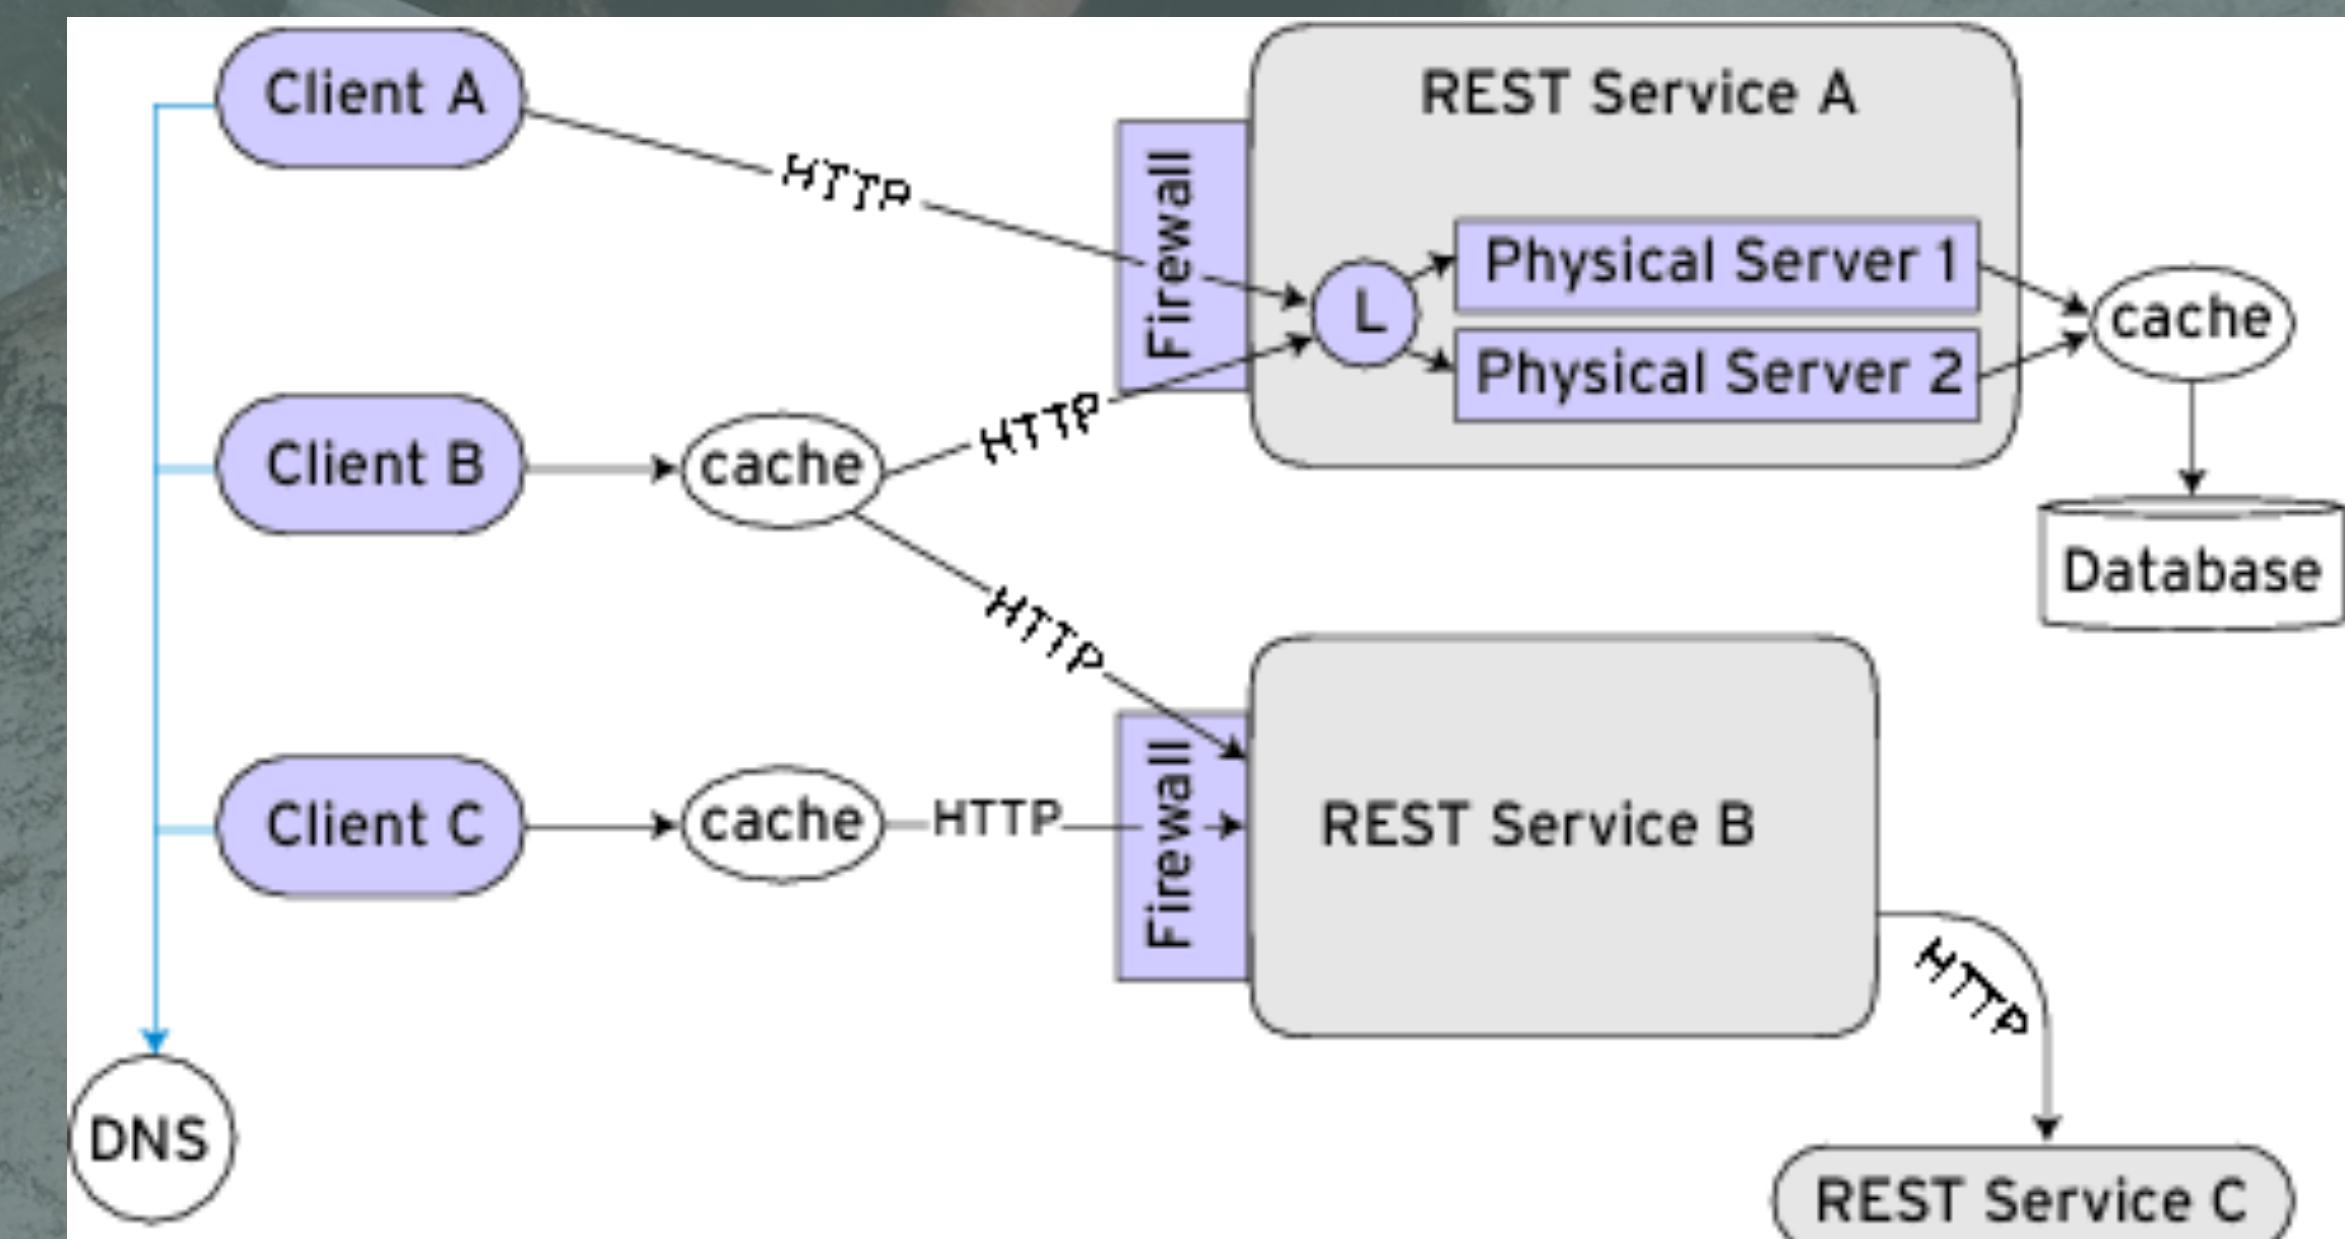


- Low latency? How low?
 - > 1 minute?
 - Consider periodic batch jobs!

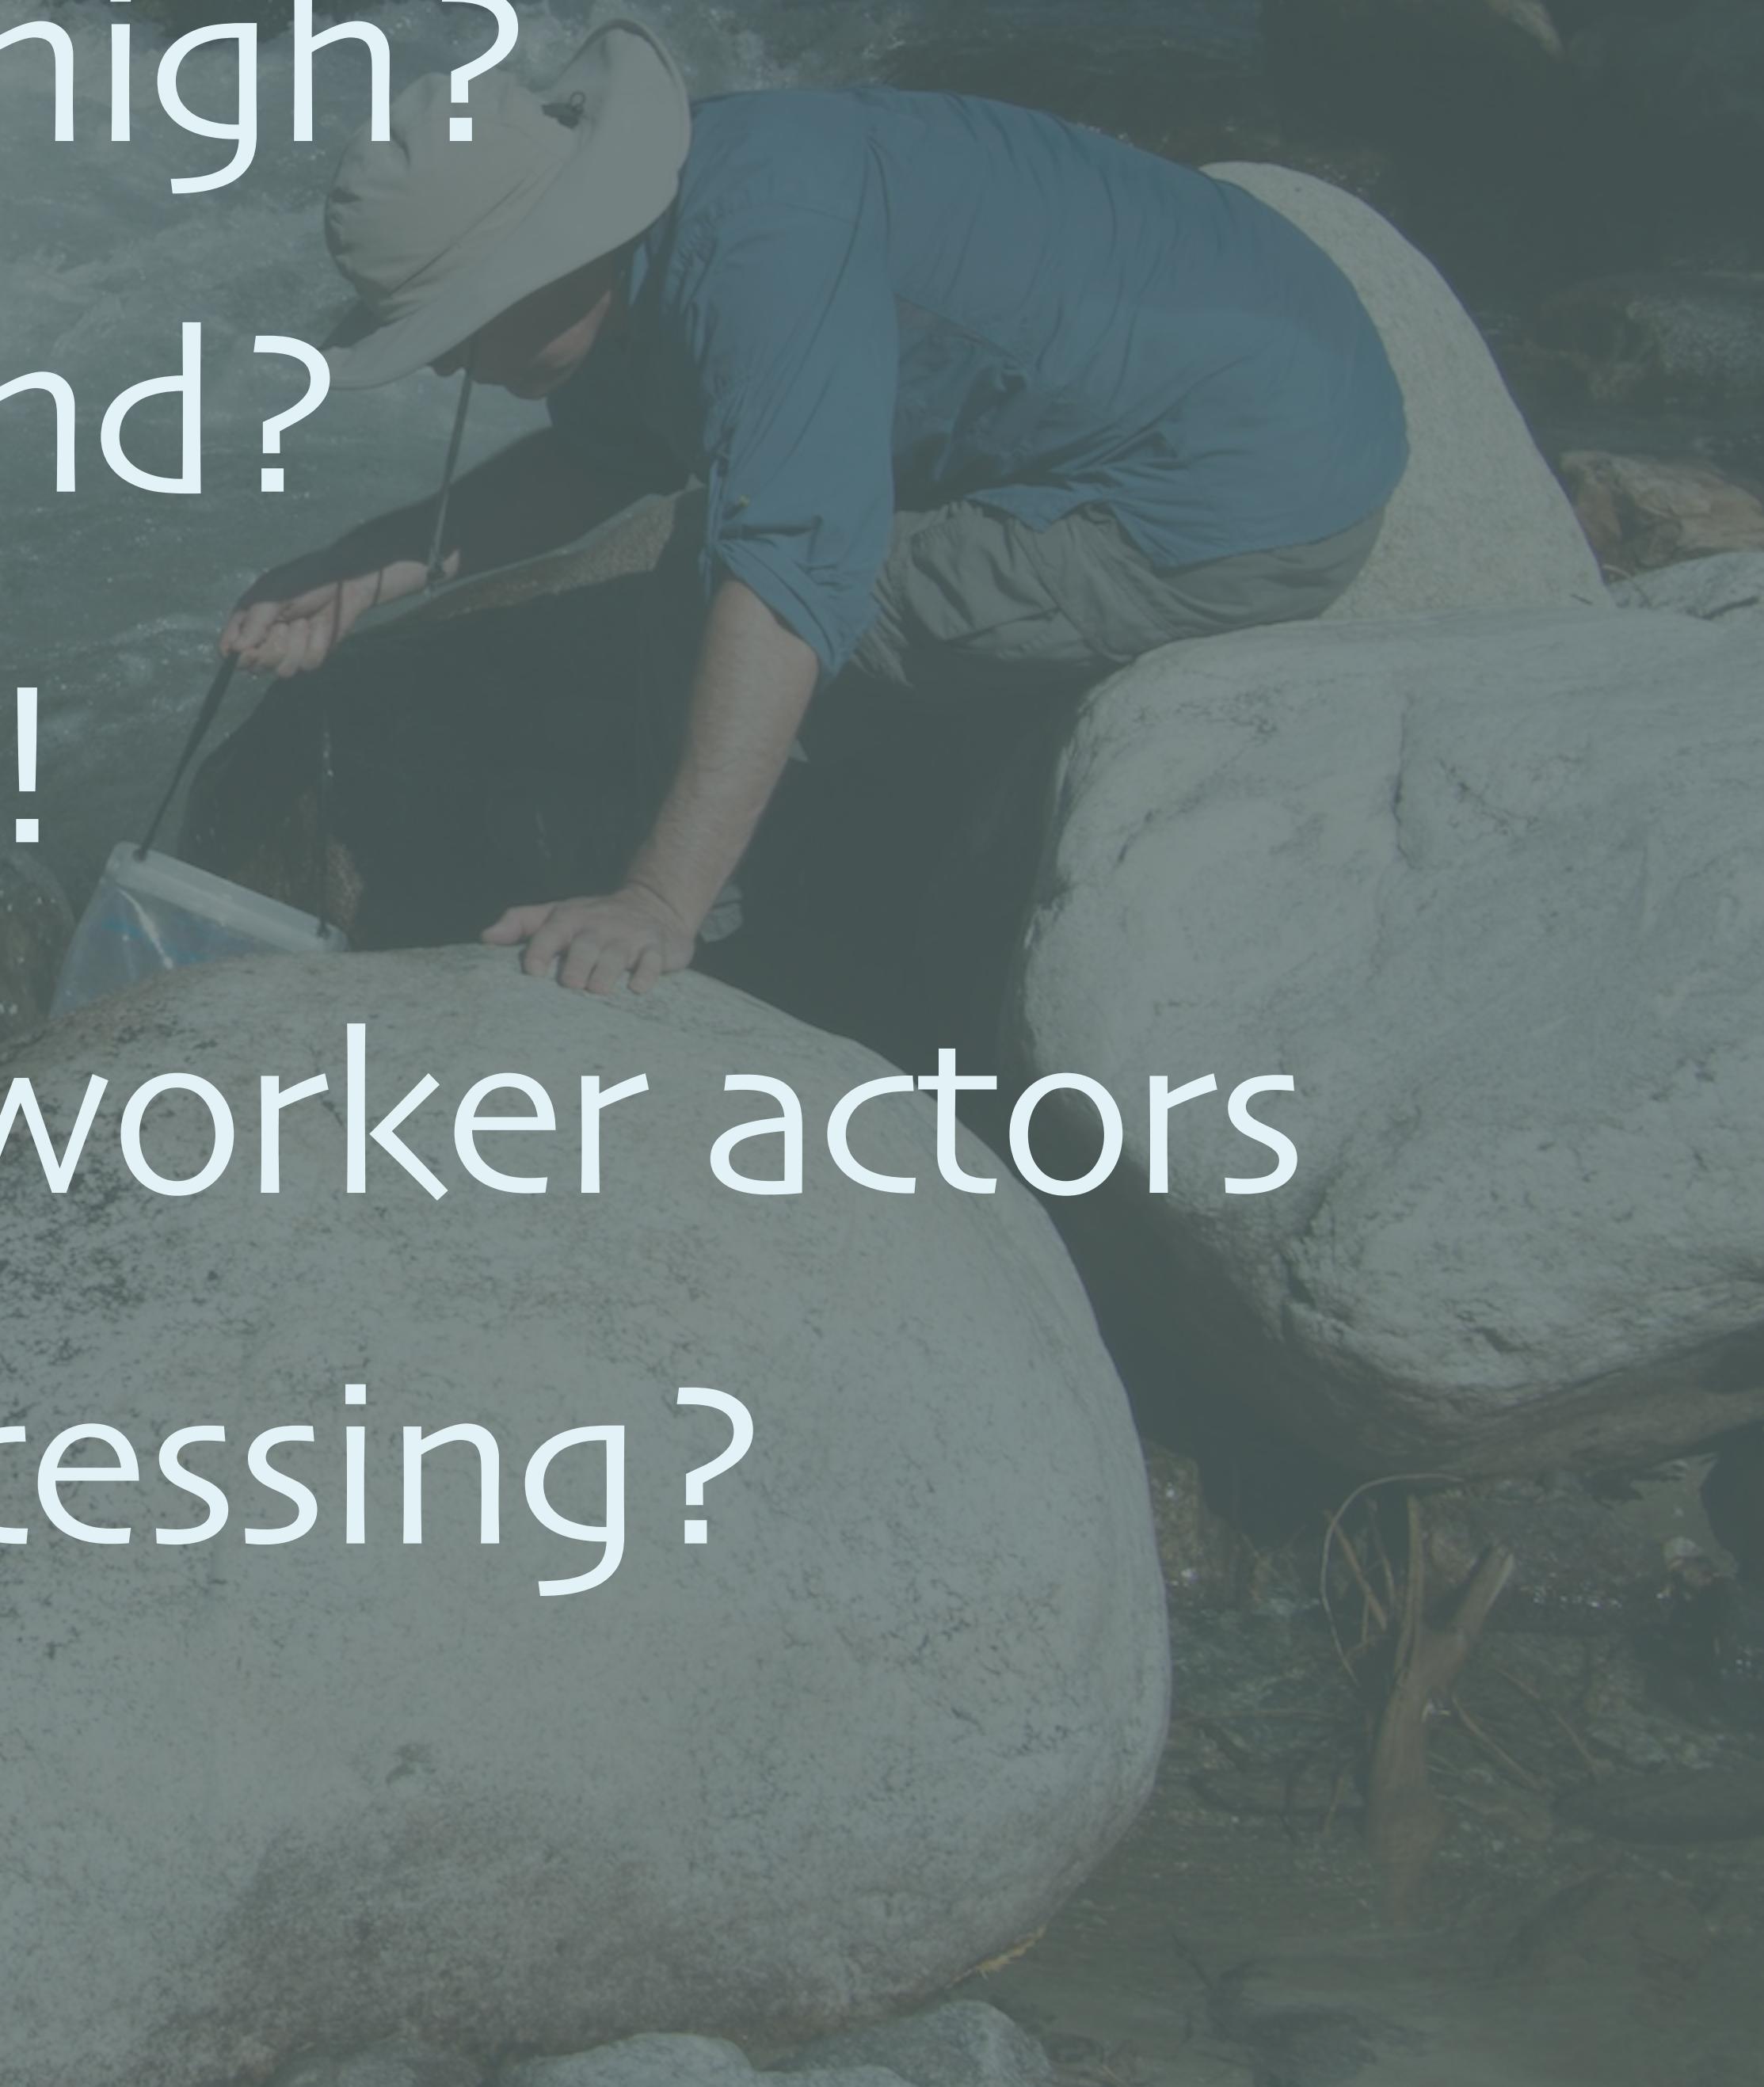
- High Volume: How high?



- High Volume: How high?
- < 10,000 events/second?
- REST
- One at a time...



- High Volume: How high?
- < 100,000 per second?
- Nonblocking REST!
- Parallelism - Akka worker actors
- Switch to bulk processing?



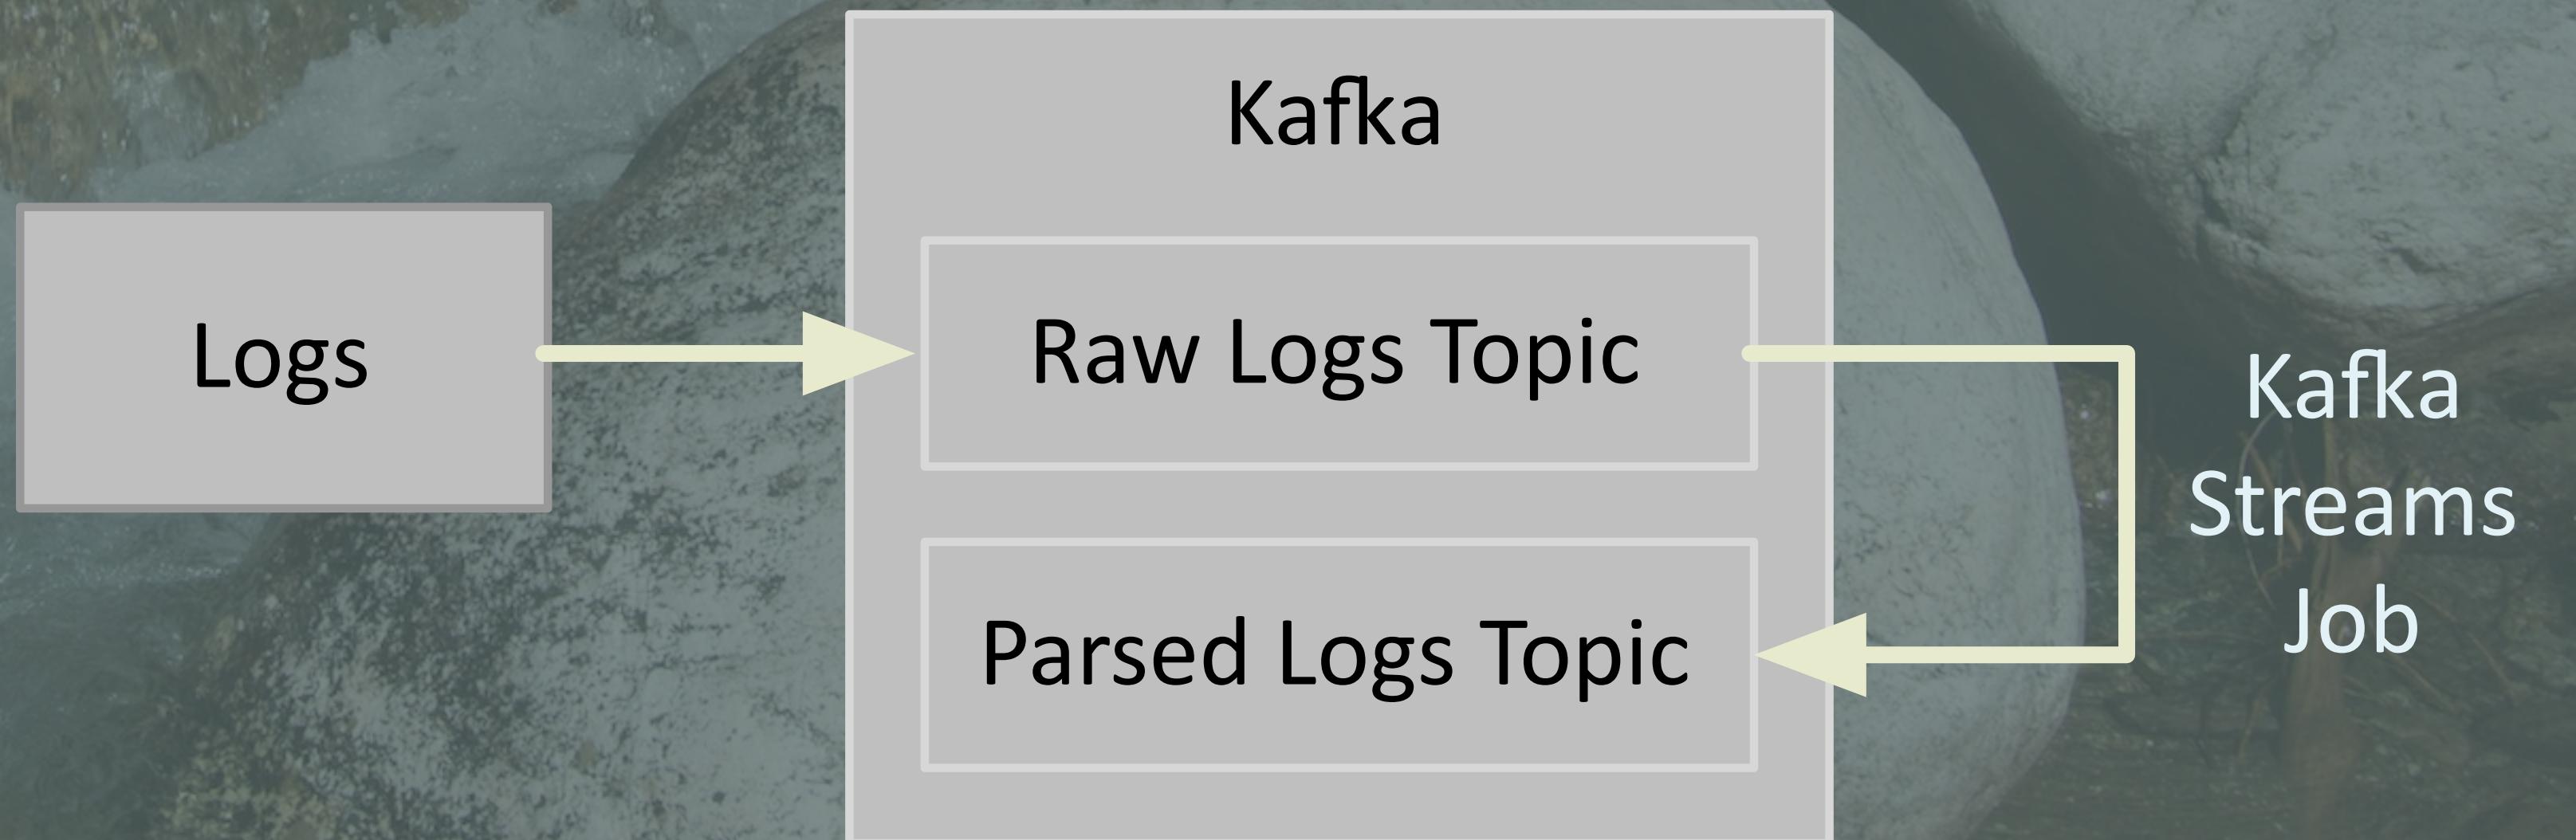
- High Volume: How high?
- 1,000,000s per second?
 - Flink or Spark Streaming
 - Process in bulk



- Which kinds of data processing?



- Which kinds of data processing?
- Extract, transform, and load (ETL)?



- Which kinds of data processing?
- “Dataflow” pipelines

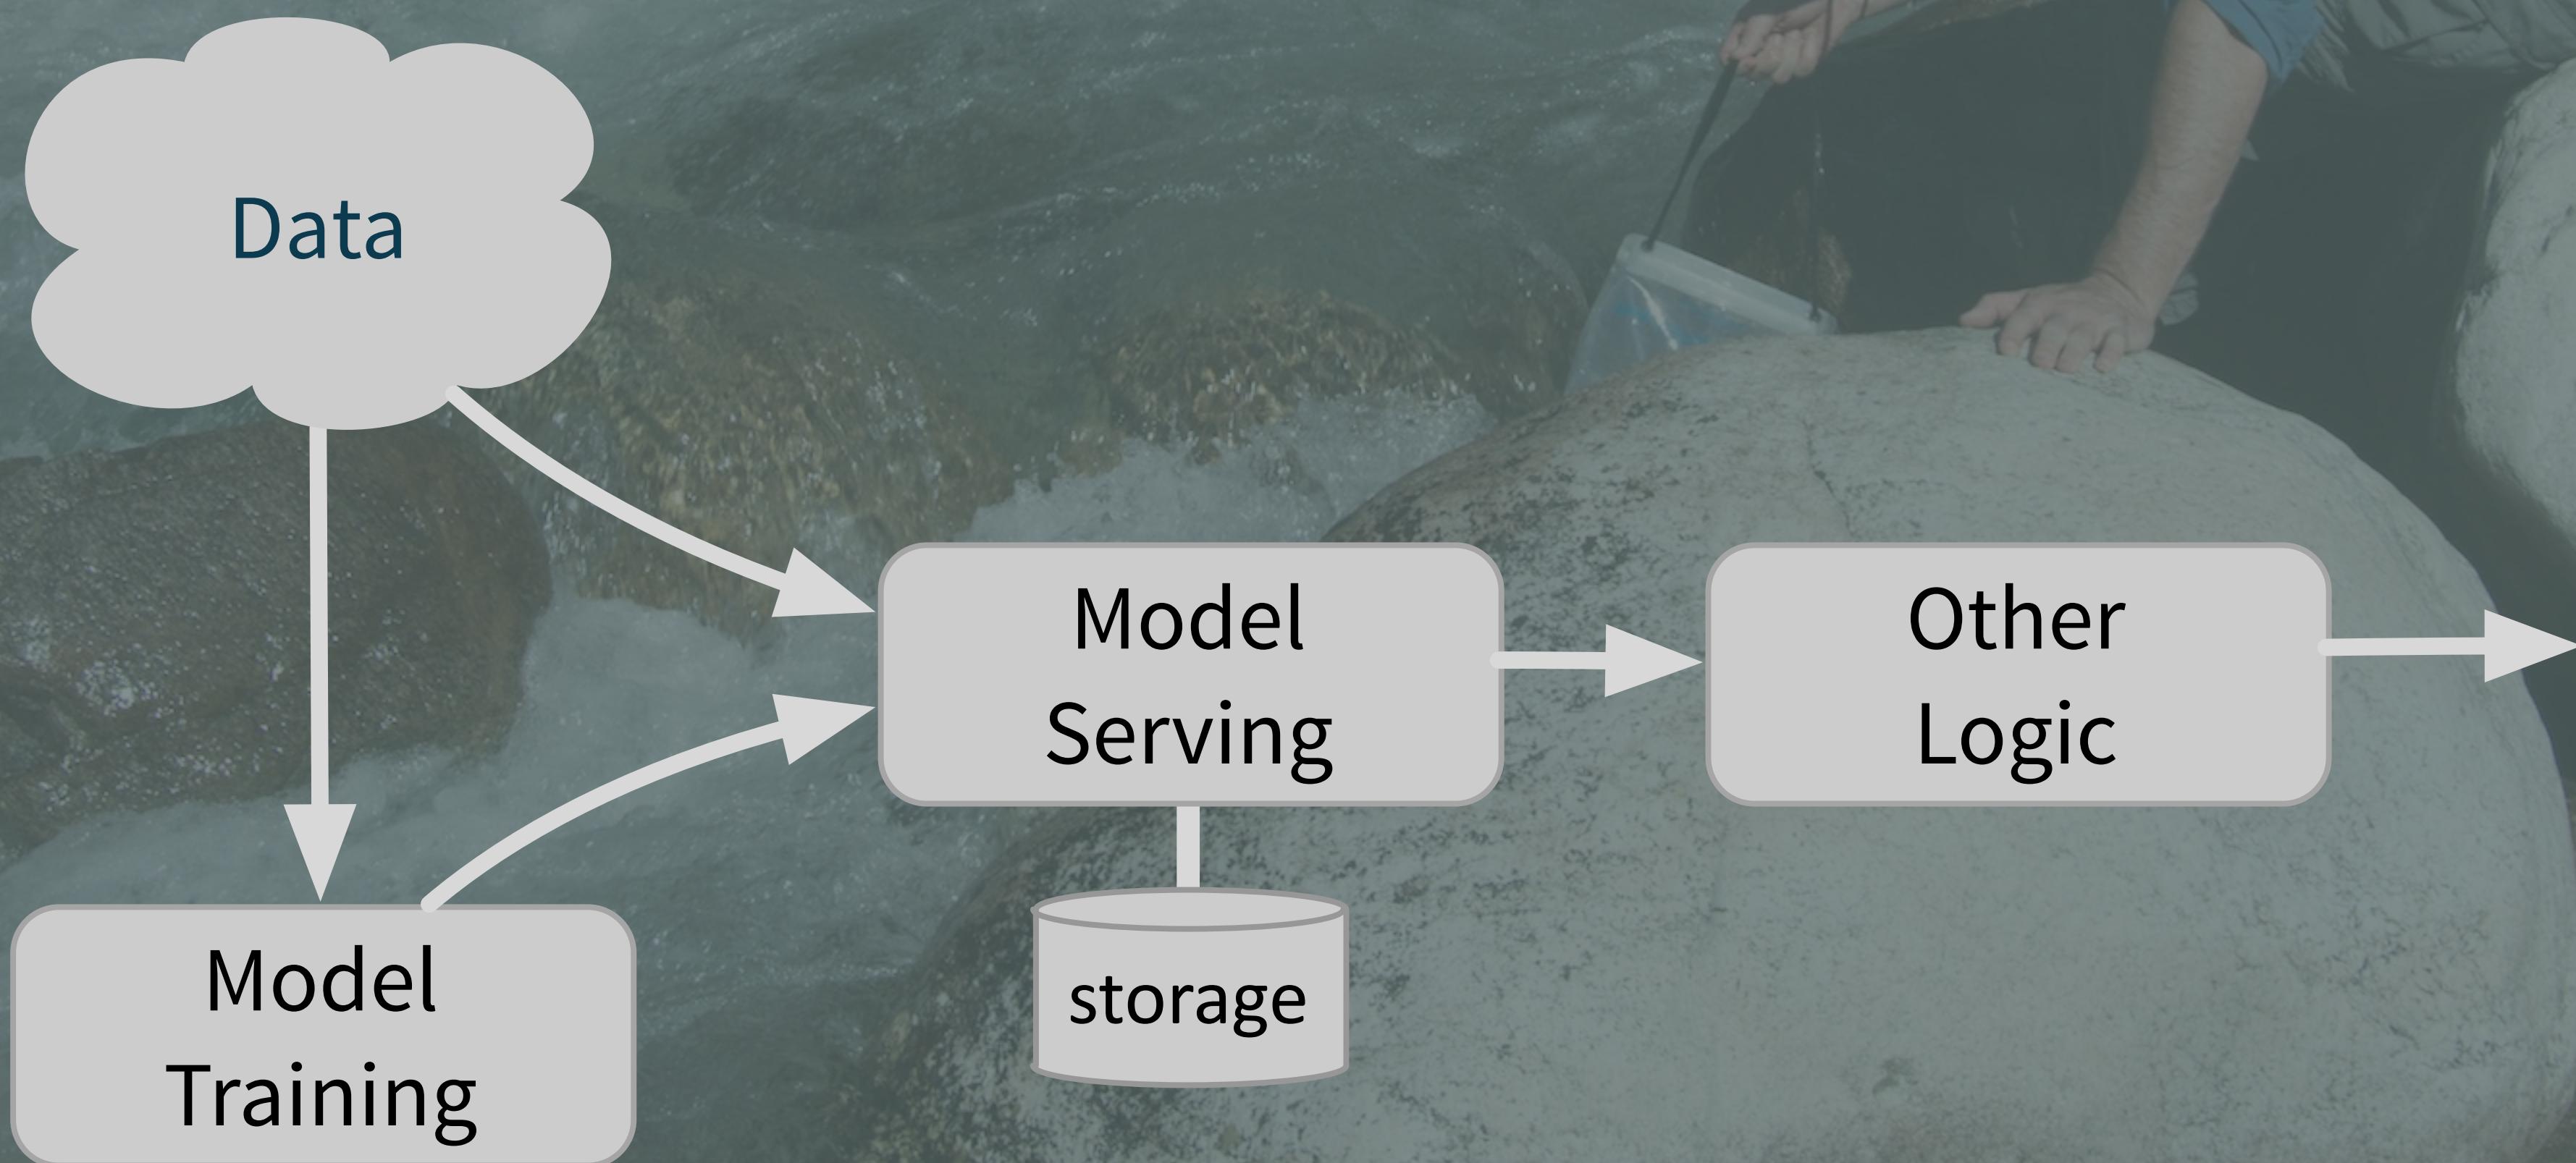
```
val sc = new SparkContext("local[*]", "Inverted Idx")
sc.textFile("data/crawl")
  .map { line => val Array(path, text) = line.split("\t",2);
(path, text)
} flatMap {
  case (path, text) => text.split("""\W+""").map((_, path))
} map {
  case (w, p) => ((w, p), 1)
} reduceByKey {
  case (n1, n2) => n1 + n2
}
```

- Which kinds of data processing?
- SQL?

```
SELECT COUNT(*)  
FROM my-iot-data  
GROUP BY zip-code
```

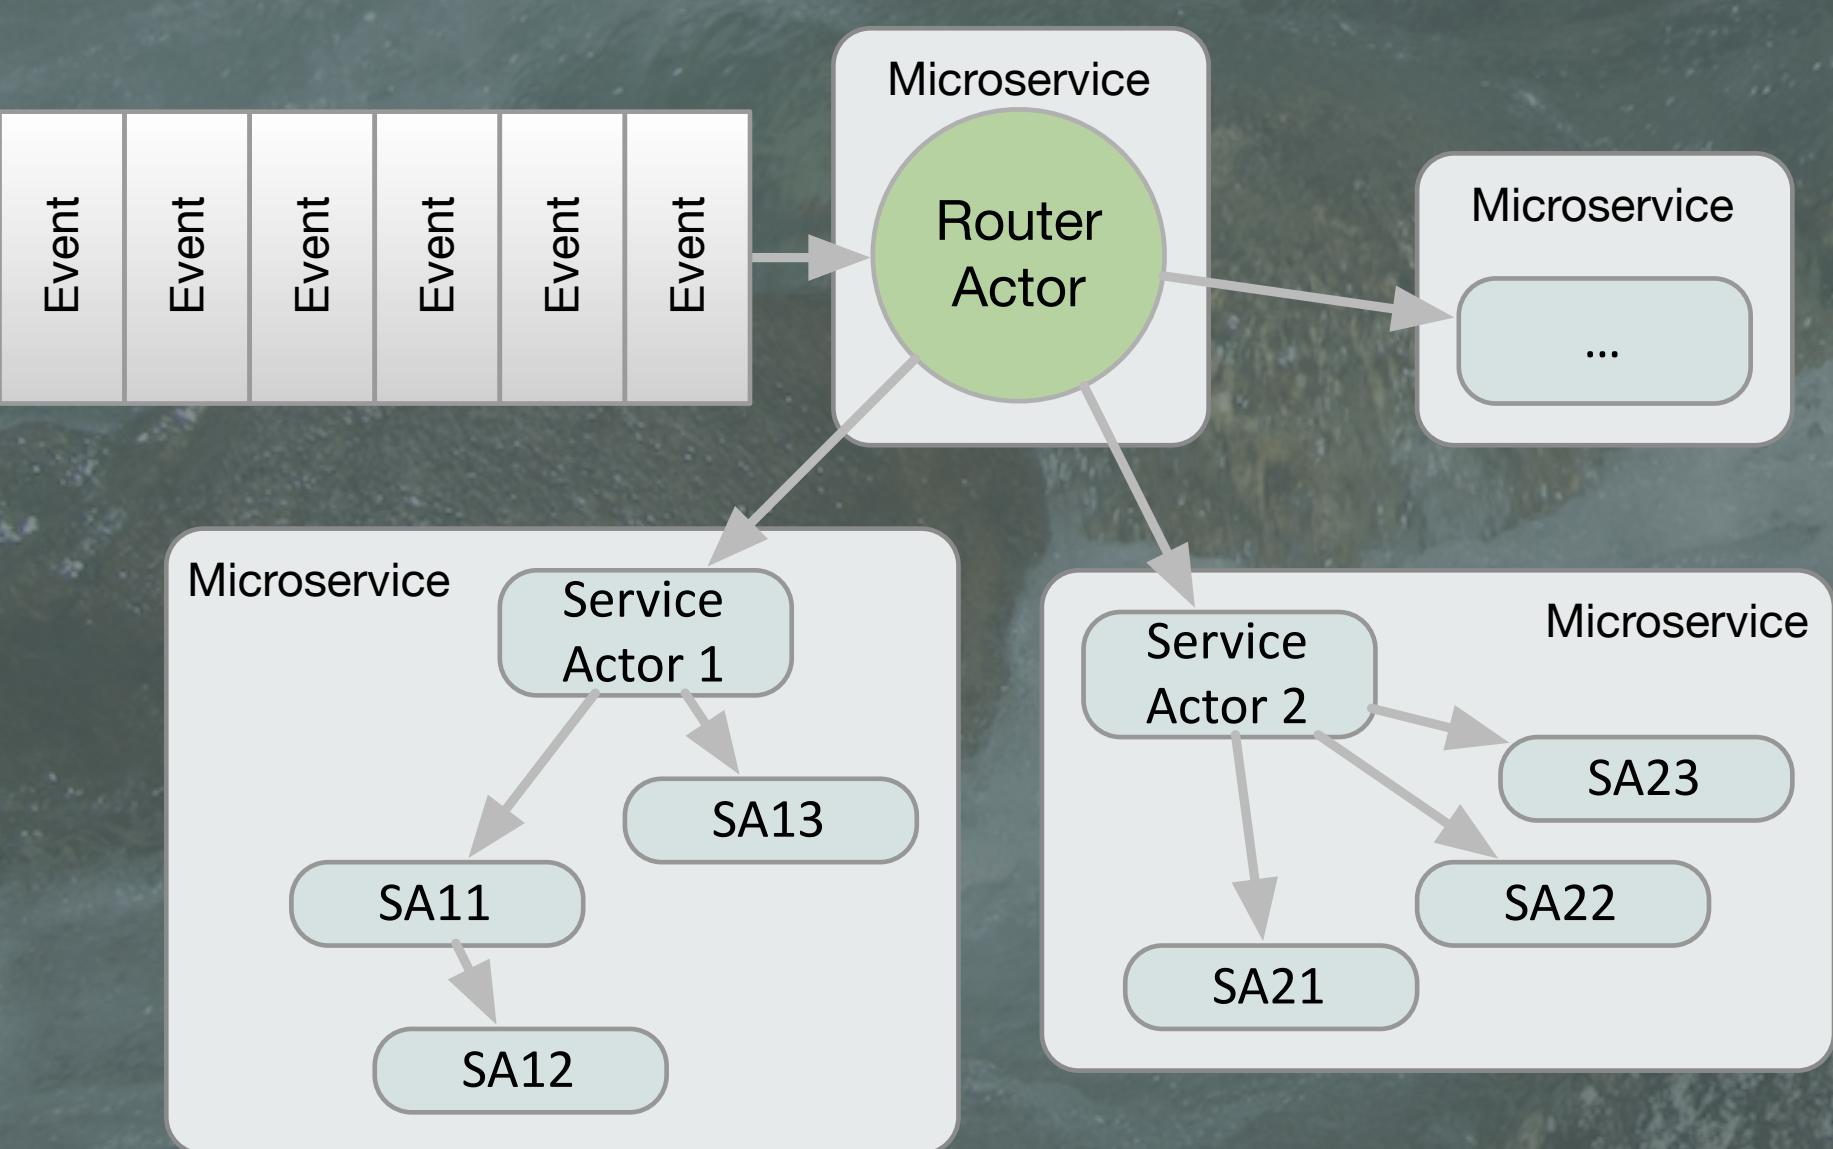
```
val input = spark.read.  
format("parquet").  
stream("my-iot-data")  
  
input.groupBy("zip-code").  
count()
```

- Which kinds of data processing?
- Train and serve ML models?



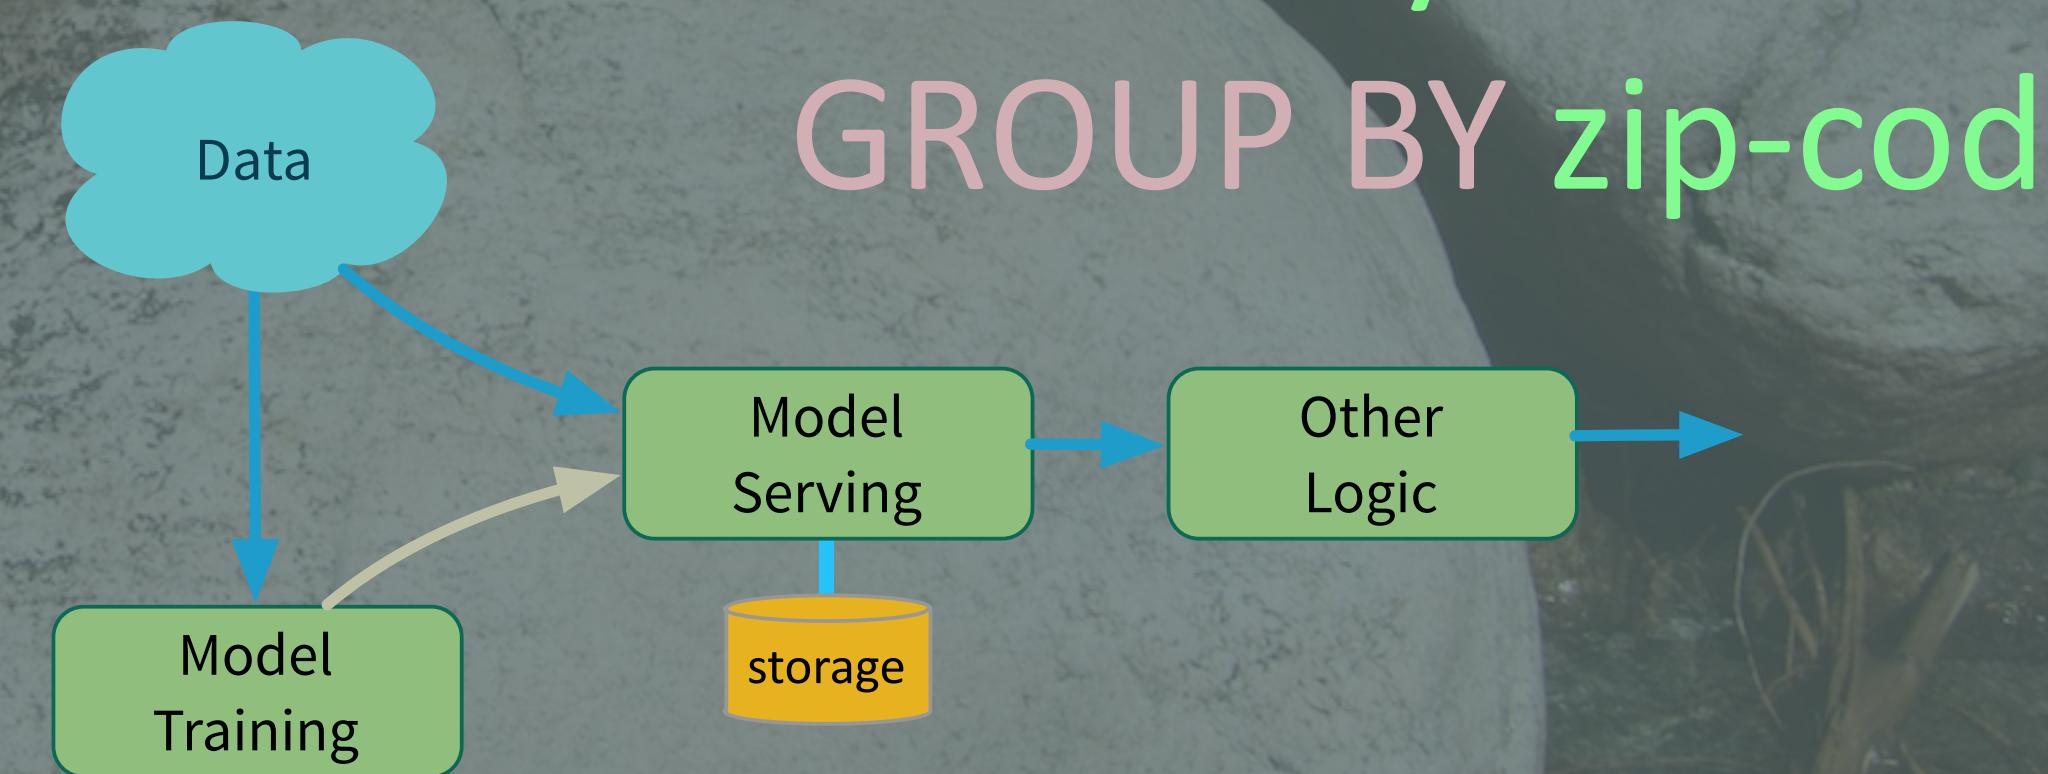
- Process data individually or in bulk?

Event-driven μ-services



"Record-centric" μ-services

`SELECT COUNT(*)
FROM my-iot-data
GROUP BY zip-code`



Events

Records

- Preferred application architecture?

• Streaming library in an app?

• or, distributed services

running your job?

Already have a microservices-based, DevOps CI/CD workflow? Stream processing with microservices may fit better into your environment!

Spark
Flink

Akka Streams
Kafka Streams

Spark

Mini-batch

Spark

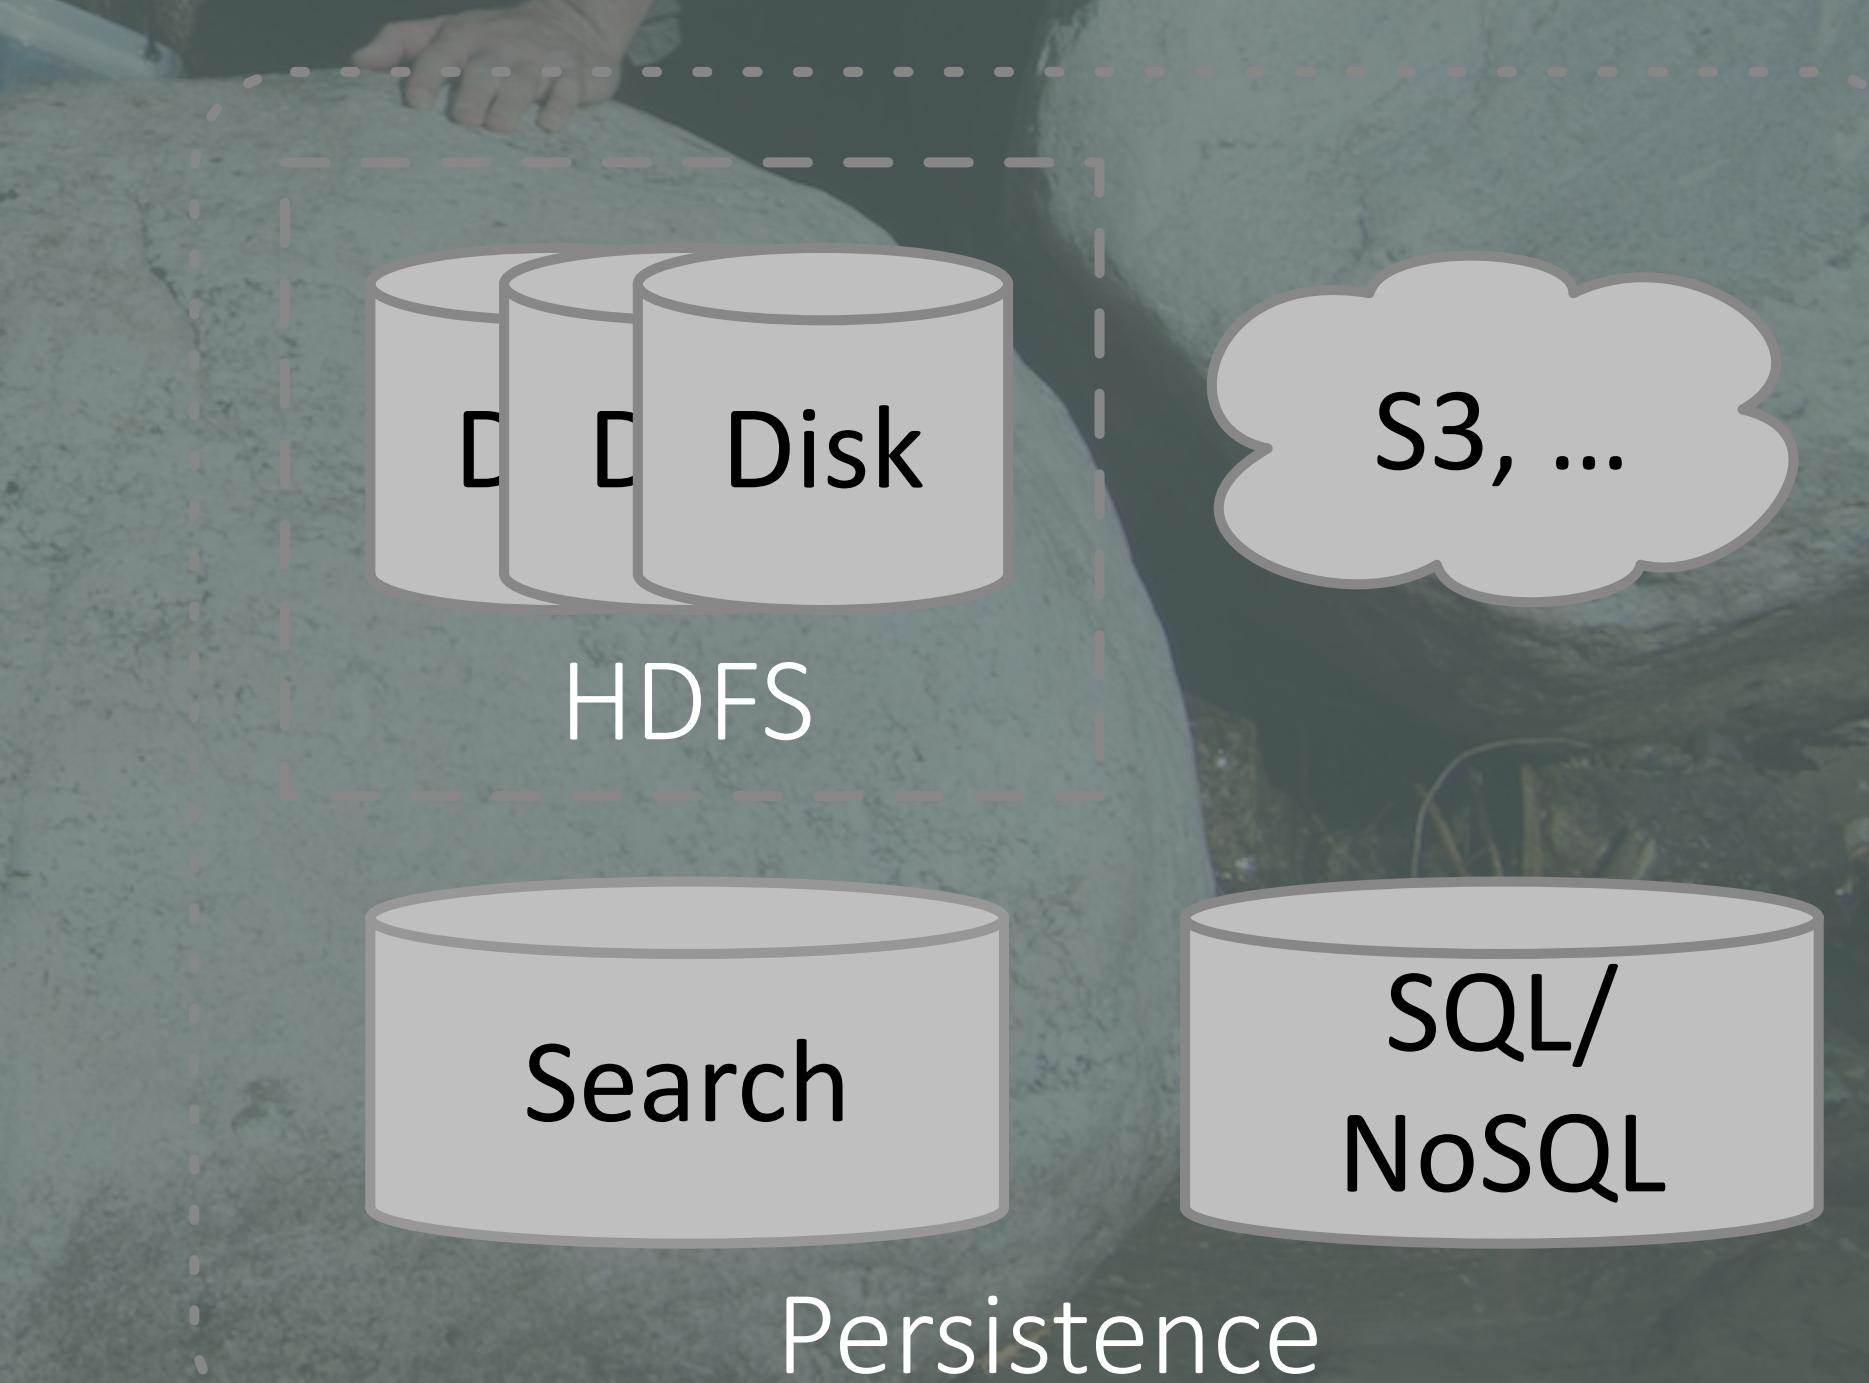
...

Batch

Beam

Low Latency

- Integration with other tools.
- Akka, Flink, & Spark integrate with Databases, Kafka, file systems, REST, ...
- Kafka Streams only read & write Kafka topics.



A close-up, low-angle shot of a waterfall cascading down a series of dark, layered rock steps. The water is white and turbulent as it falls and splashes onto the rocks below. The lighting is bright, highlighting the spray and the texture of the rocks.

Best of Breed Streaming Engines



A dark, grainy photograph of a mountainous landscape with snow-capped peaks and rocky terrain.

Spark

Flink

Beam

Akka Streams

Kafka Streams

Low Latency

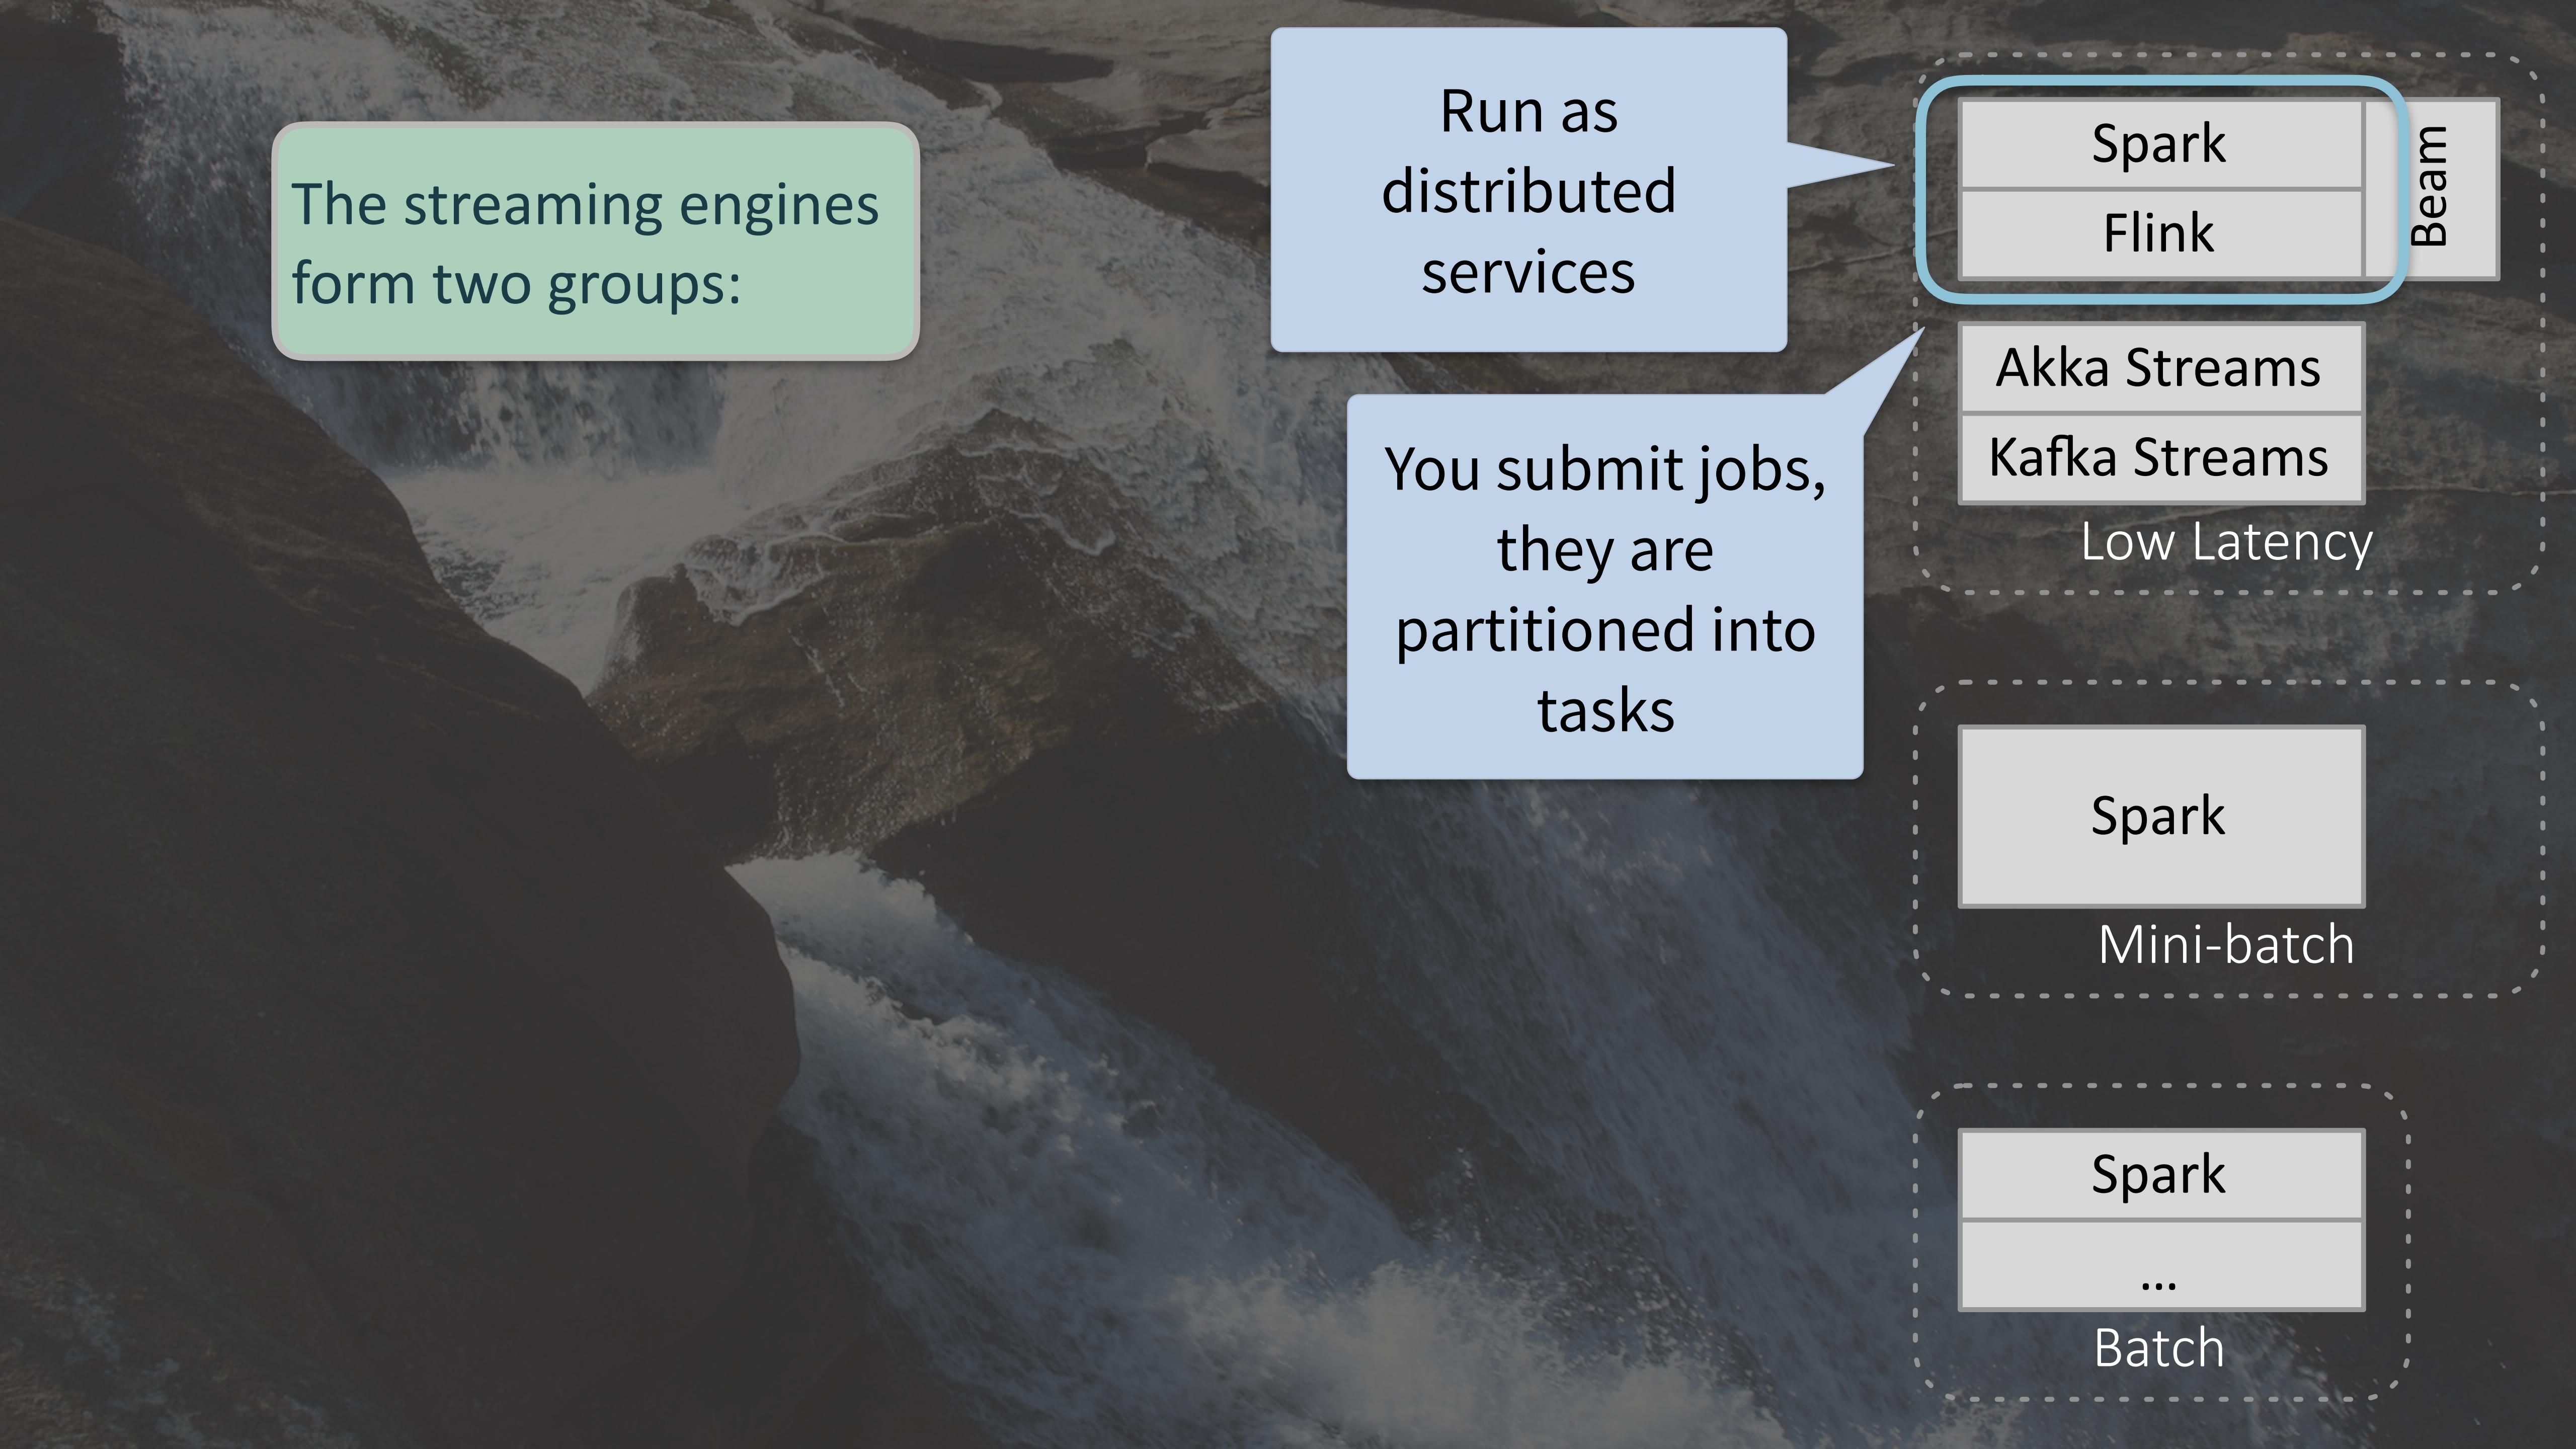
Spark

Mini-batch

Spark

...

Batch



The streaming engines form two groups:

Run as distributed services

You submit jobs, they are partitioned into tasks

Spark

Flink

Beam

Akka Streams

Kafka Streams

Low Latency

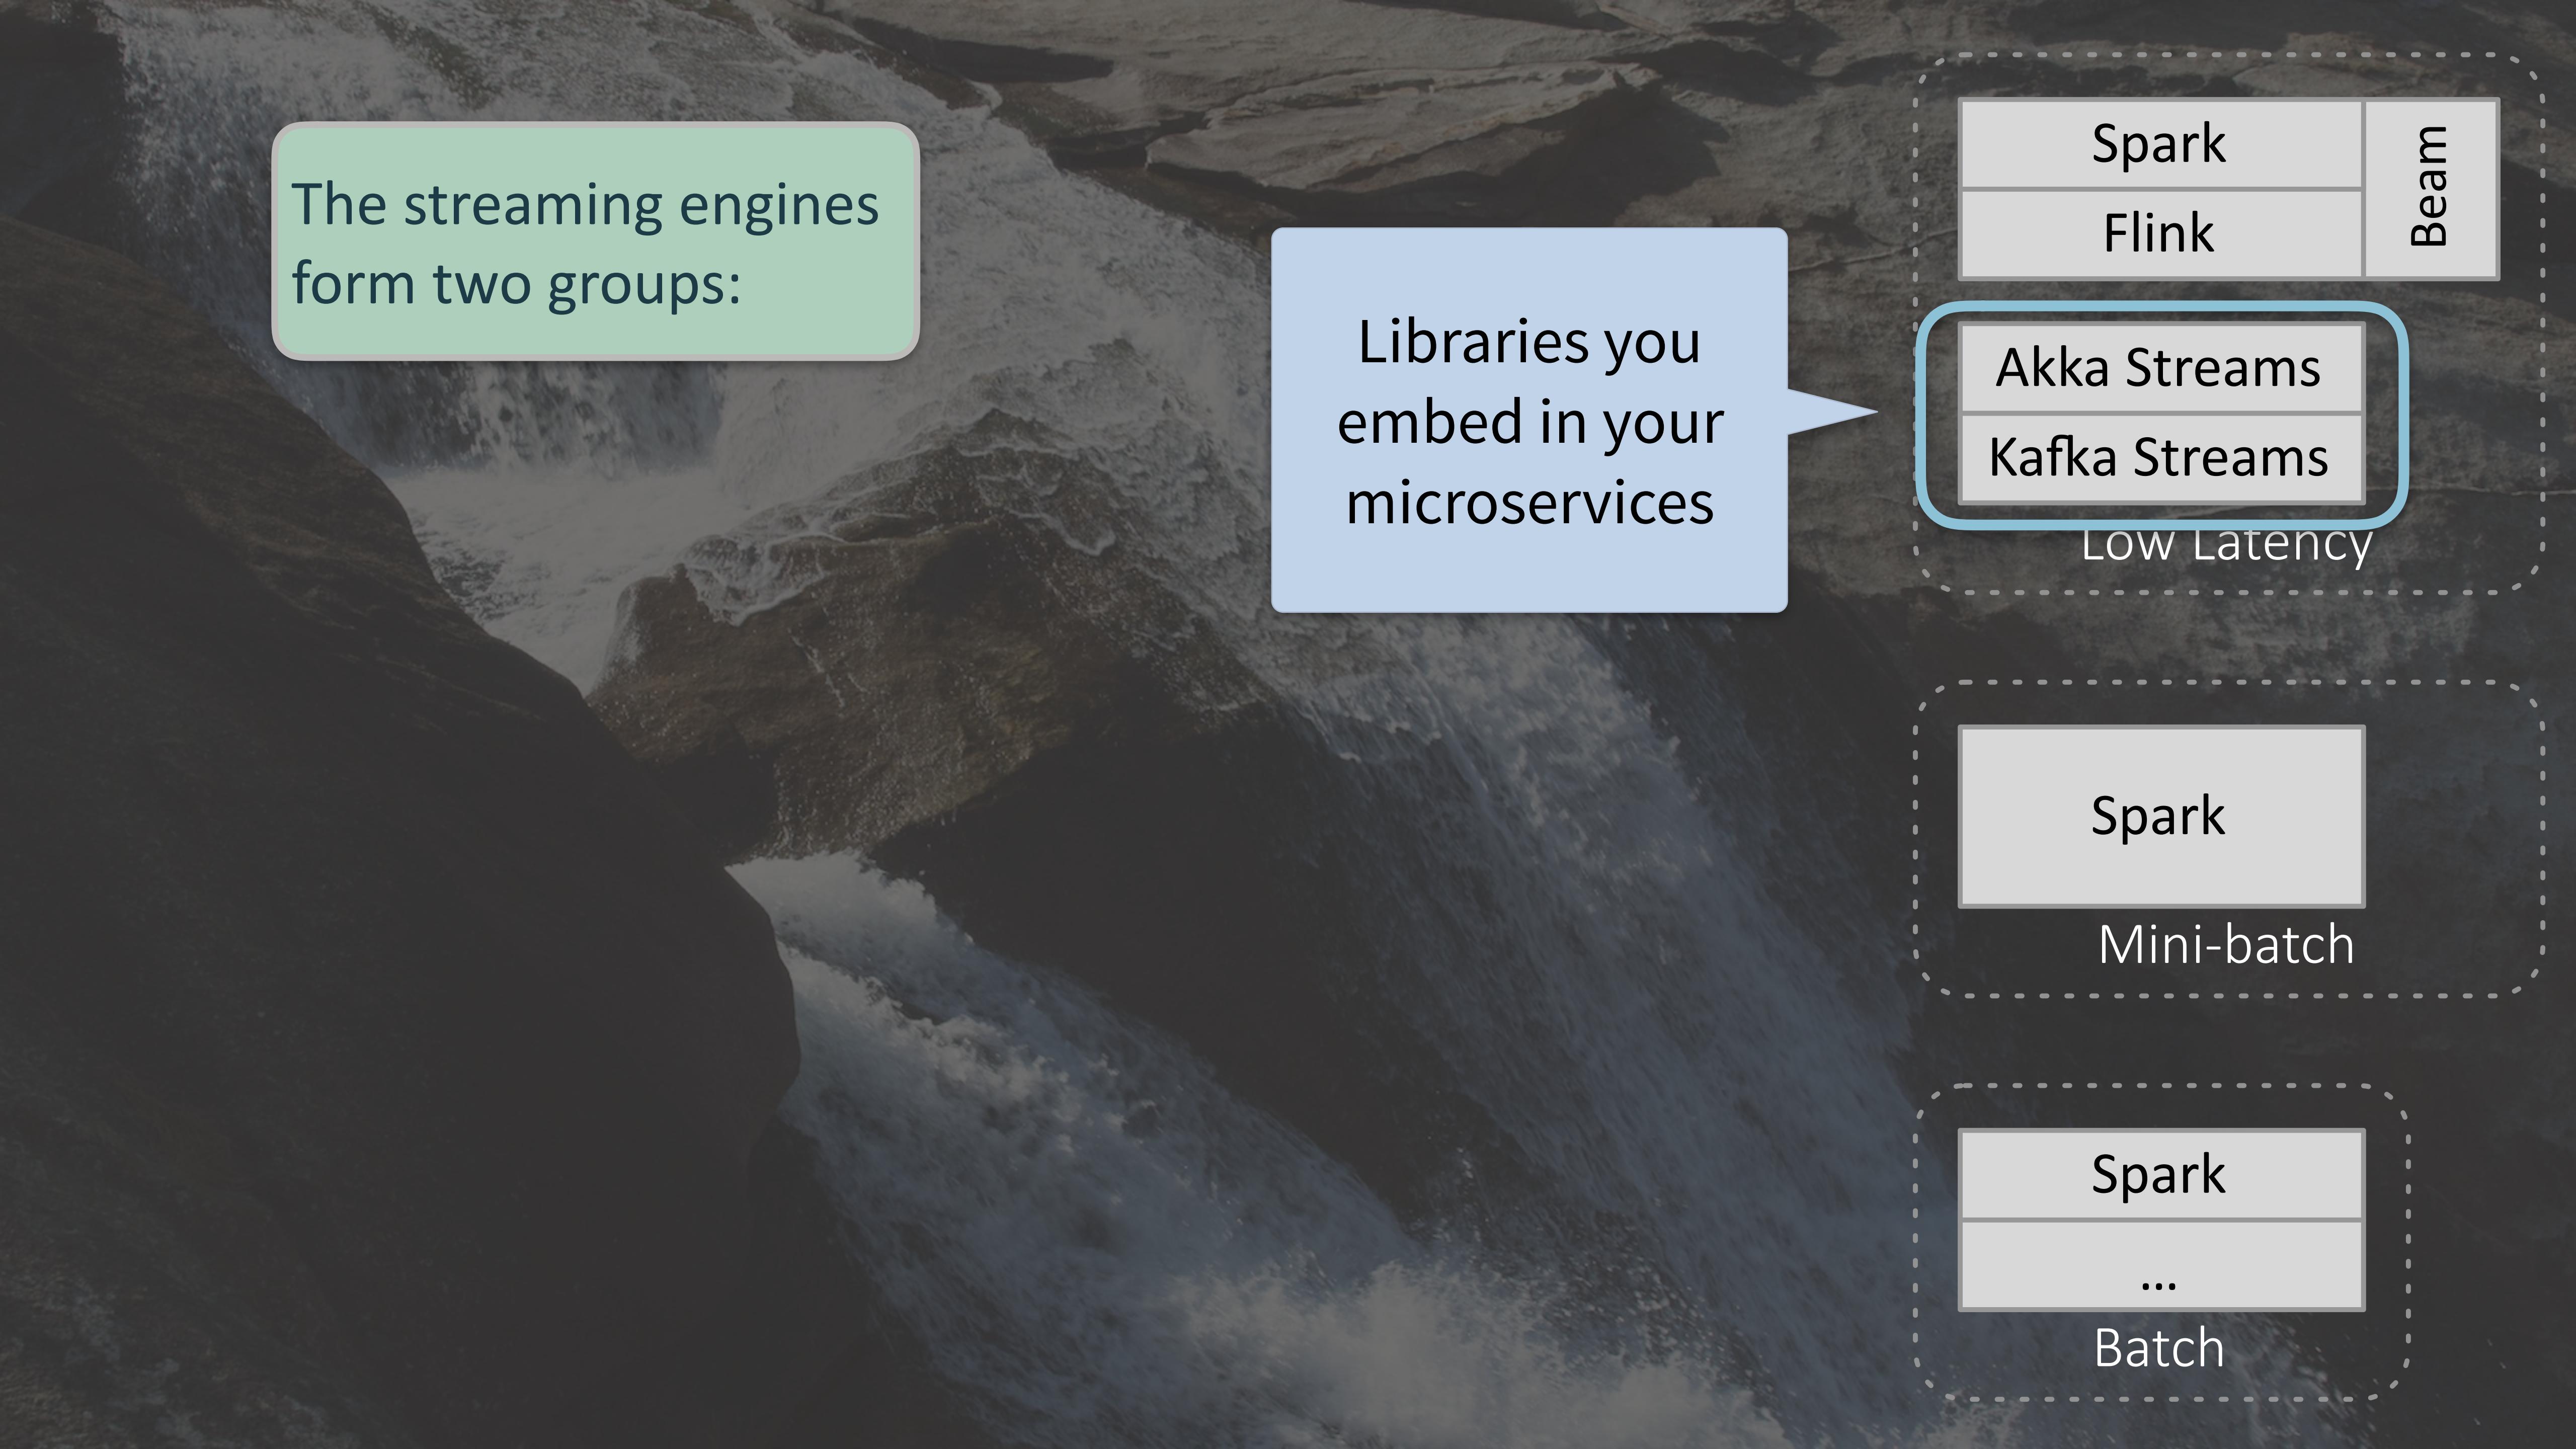
Spark

Mini-batch

Spark

...

Batch



The streaming engines form two groups:

Libraries you embed in your microservices

Spark

Flink

Beam

Akka Streams

Kafka Streams

Low Latency

Spark

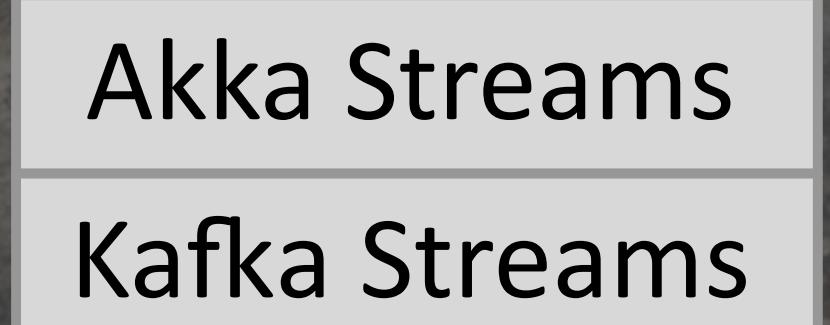
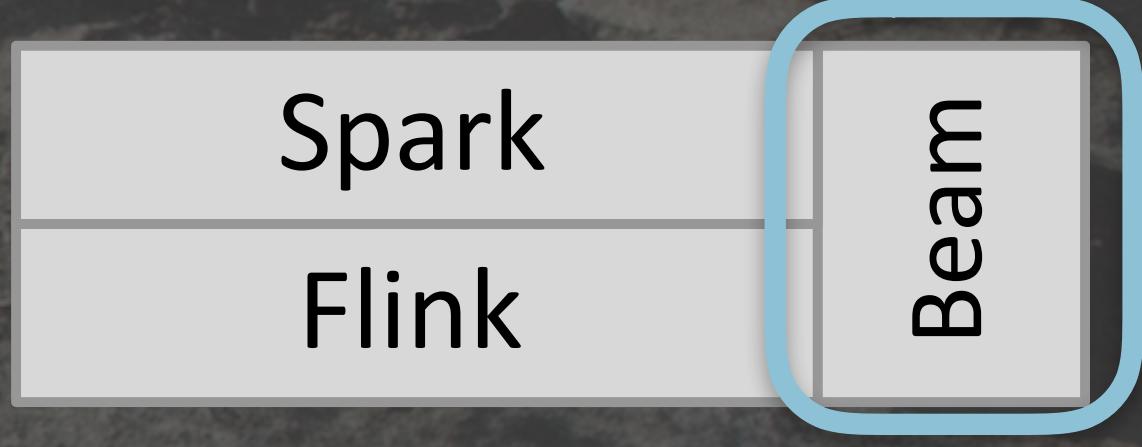
Mini-batch

Spark

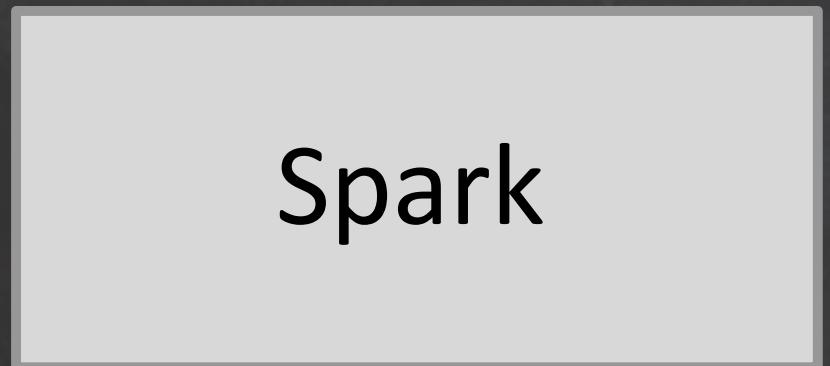
...

Batch

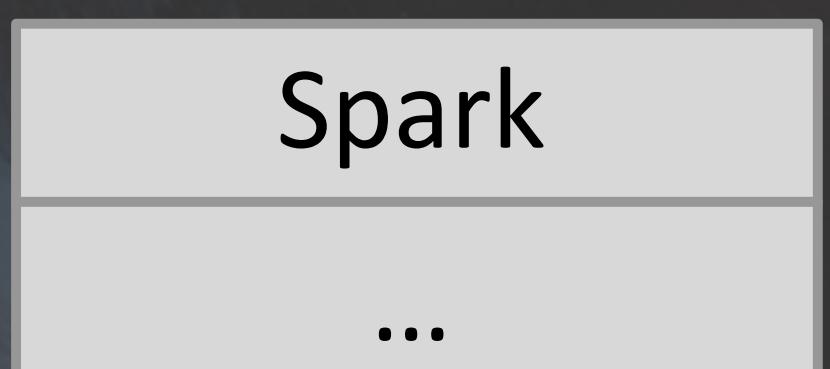
- Apache Beam
- (Google Dataflow)
- Requires a “runner”
- Most sophisticated streaming semantics



Low Latency



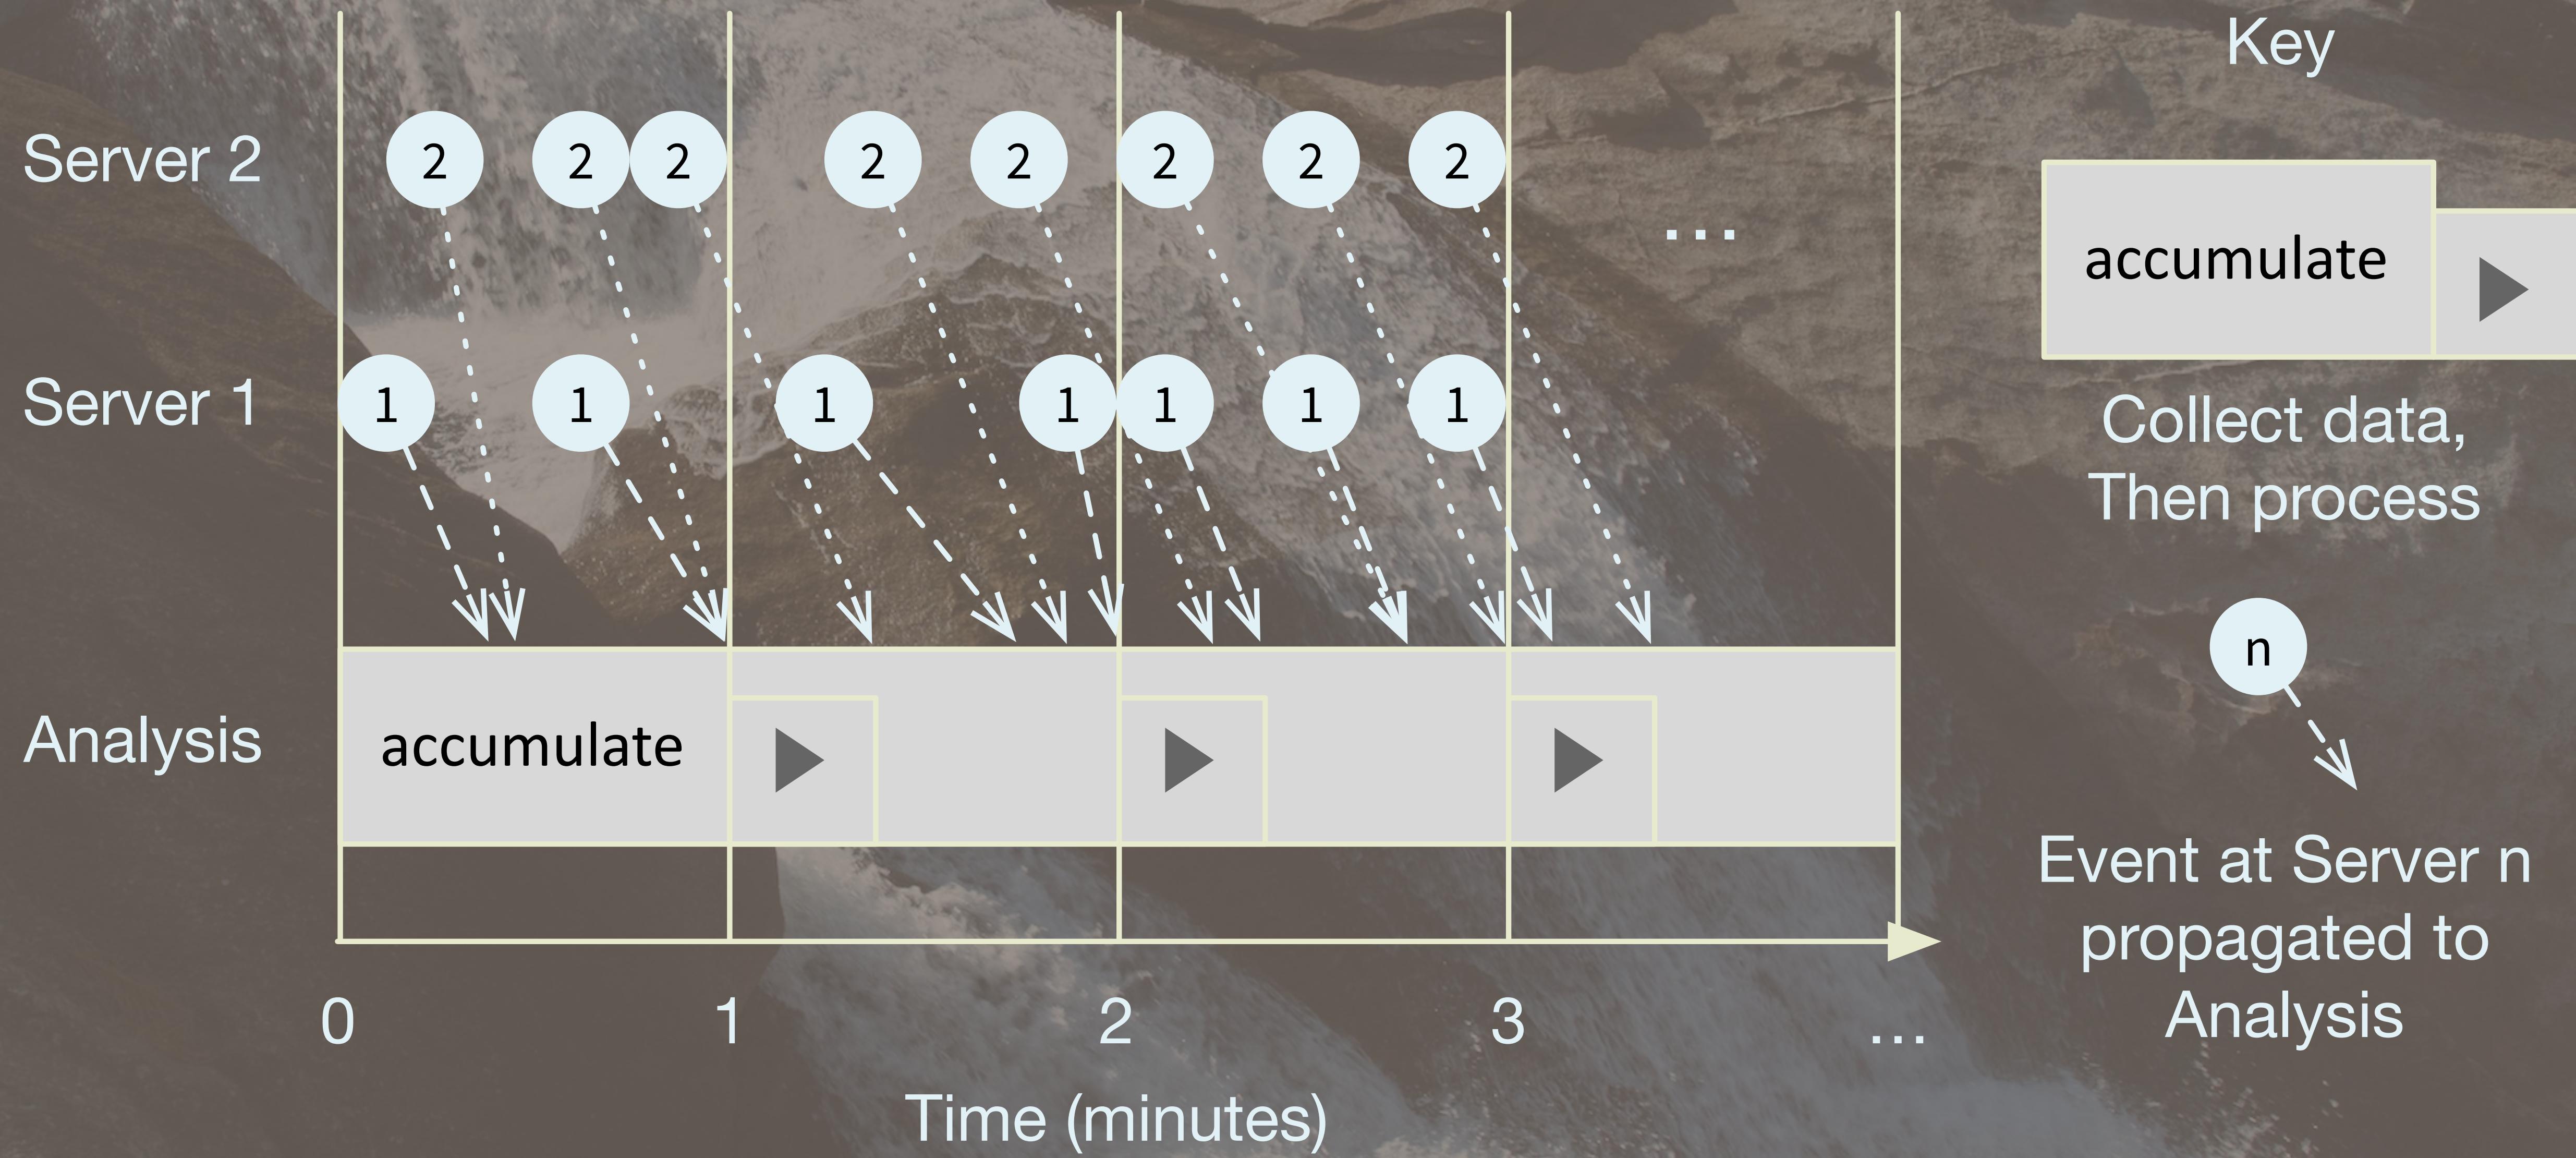
Mini-batch



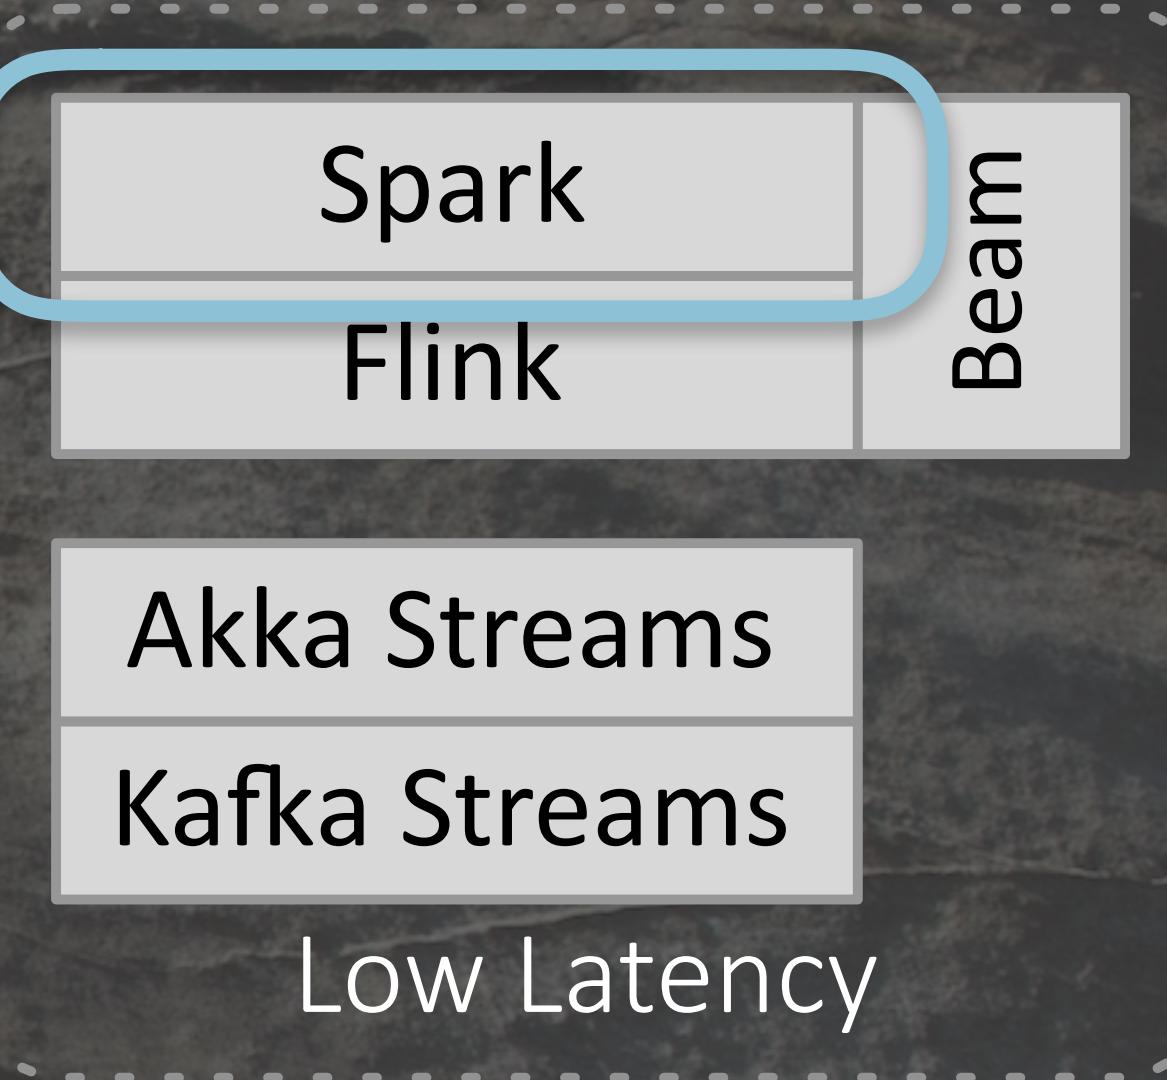
Batch
...



See these blog posts: <https://www.oreilly.com/people/09f01-tyler-akidau>



- Spark Structured Streaming
 - “Dataset” - SQL
 - Millisecond latency
 - Ideal for Rich SQL, ML.

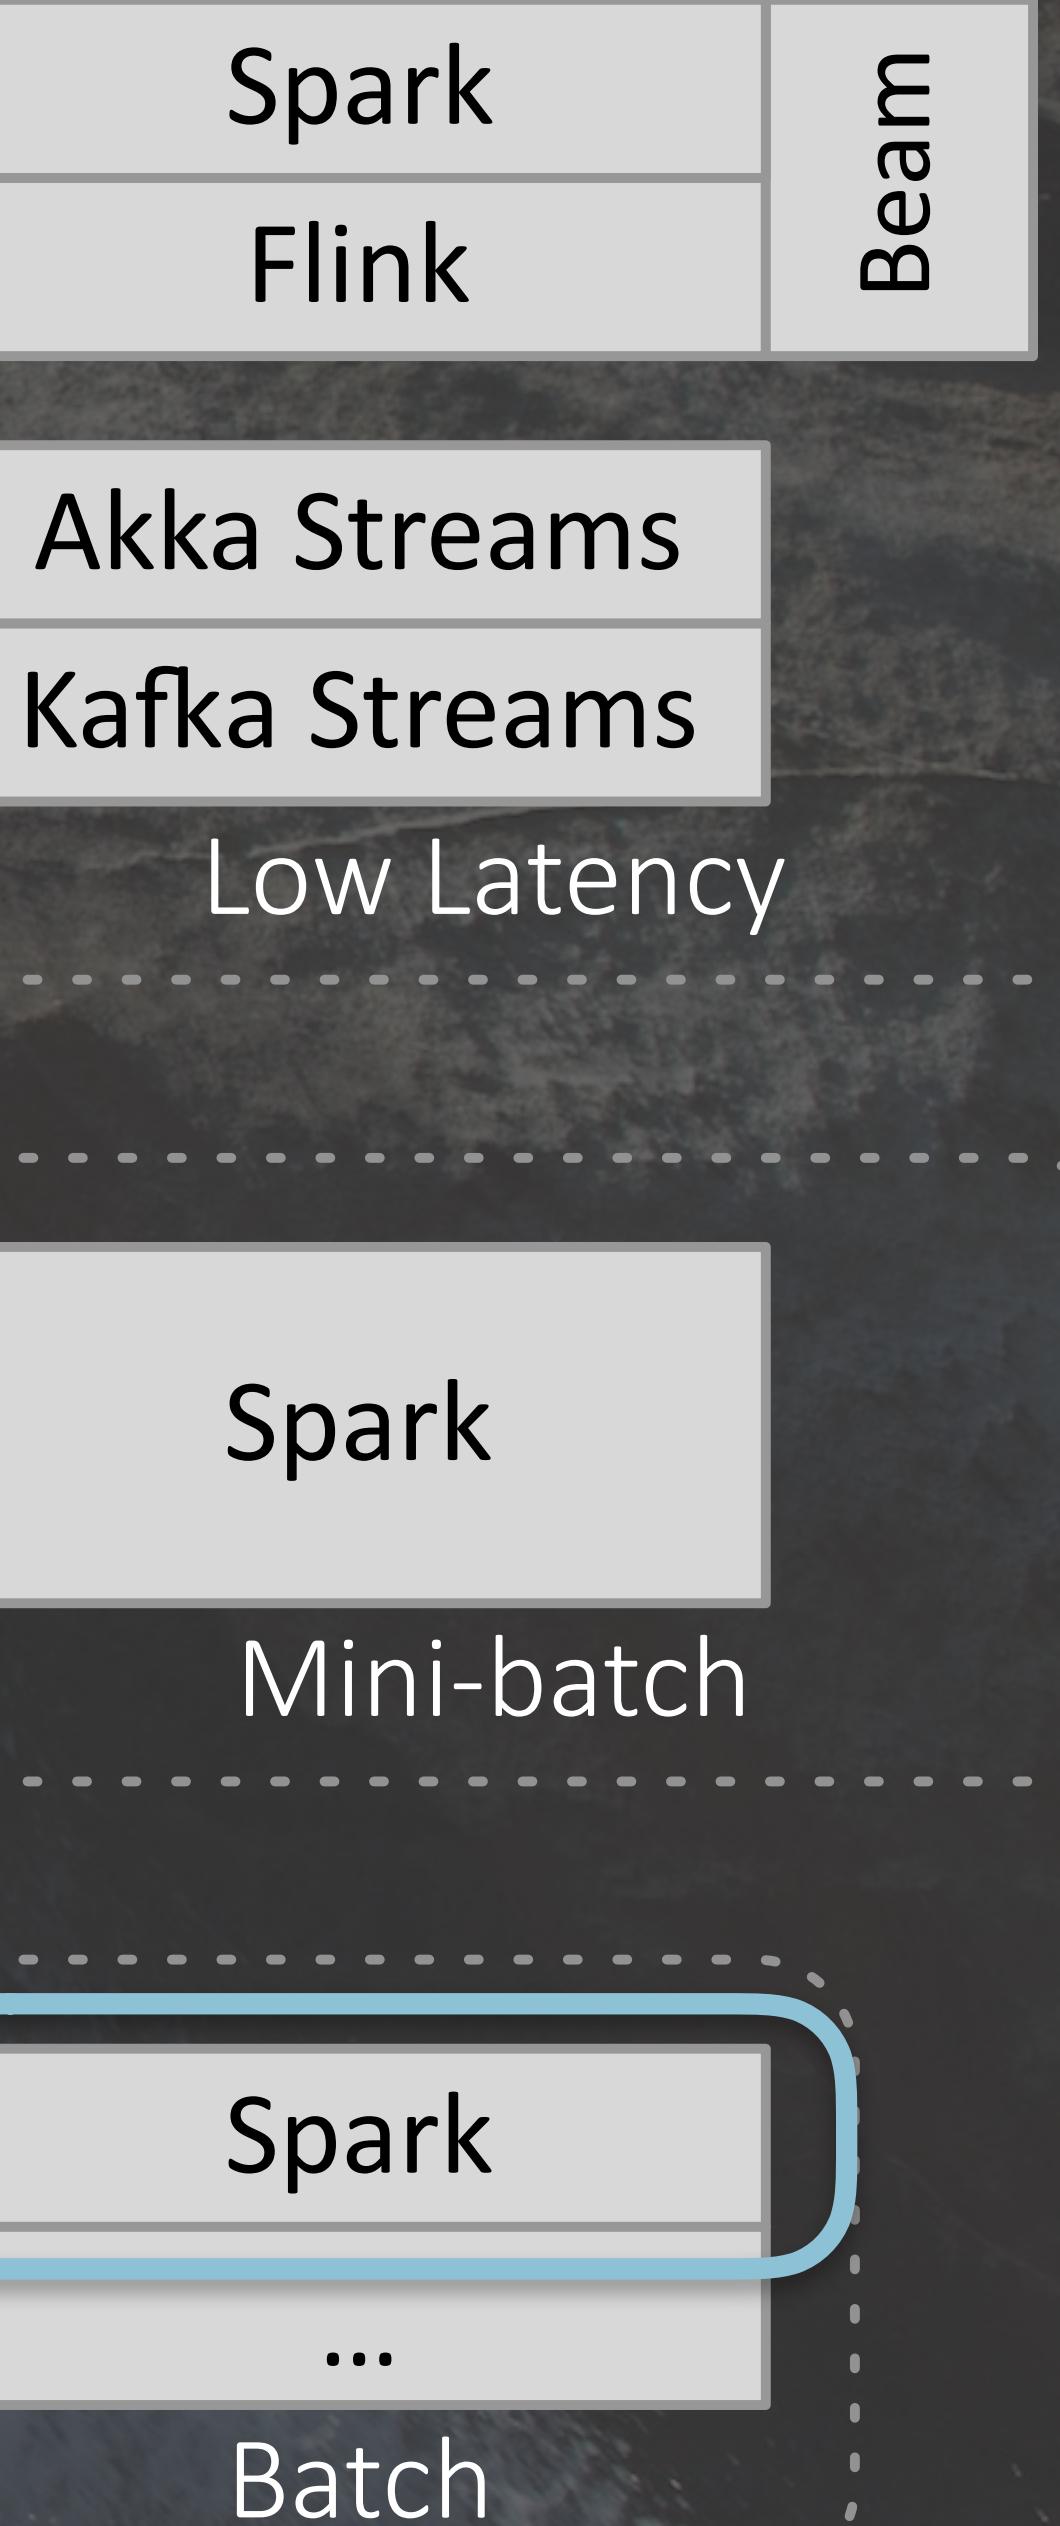


- Spark Streaming
 - Mini-batch model
 - “RDD” (dataflow) based
 - ~0.5 sec latency
- Original model - obsolete



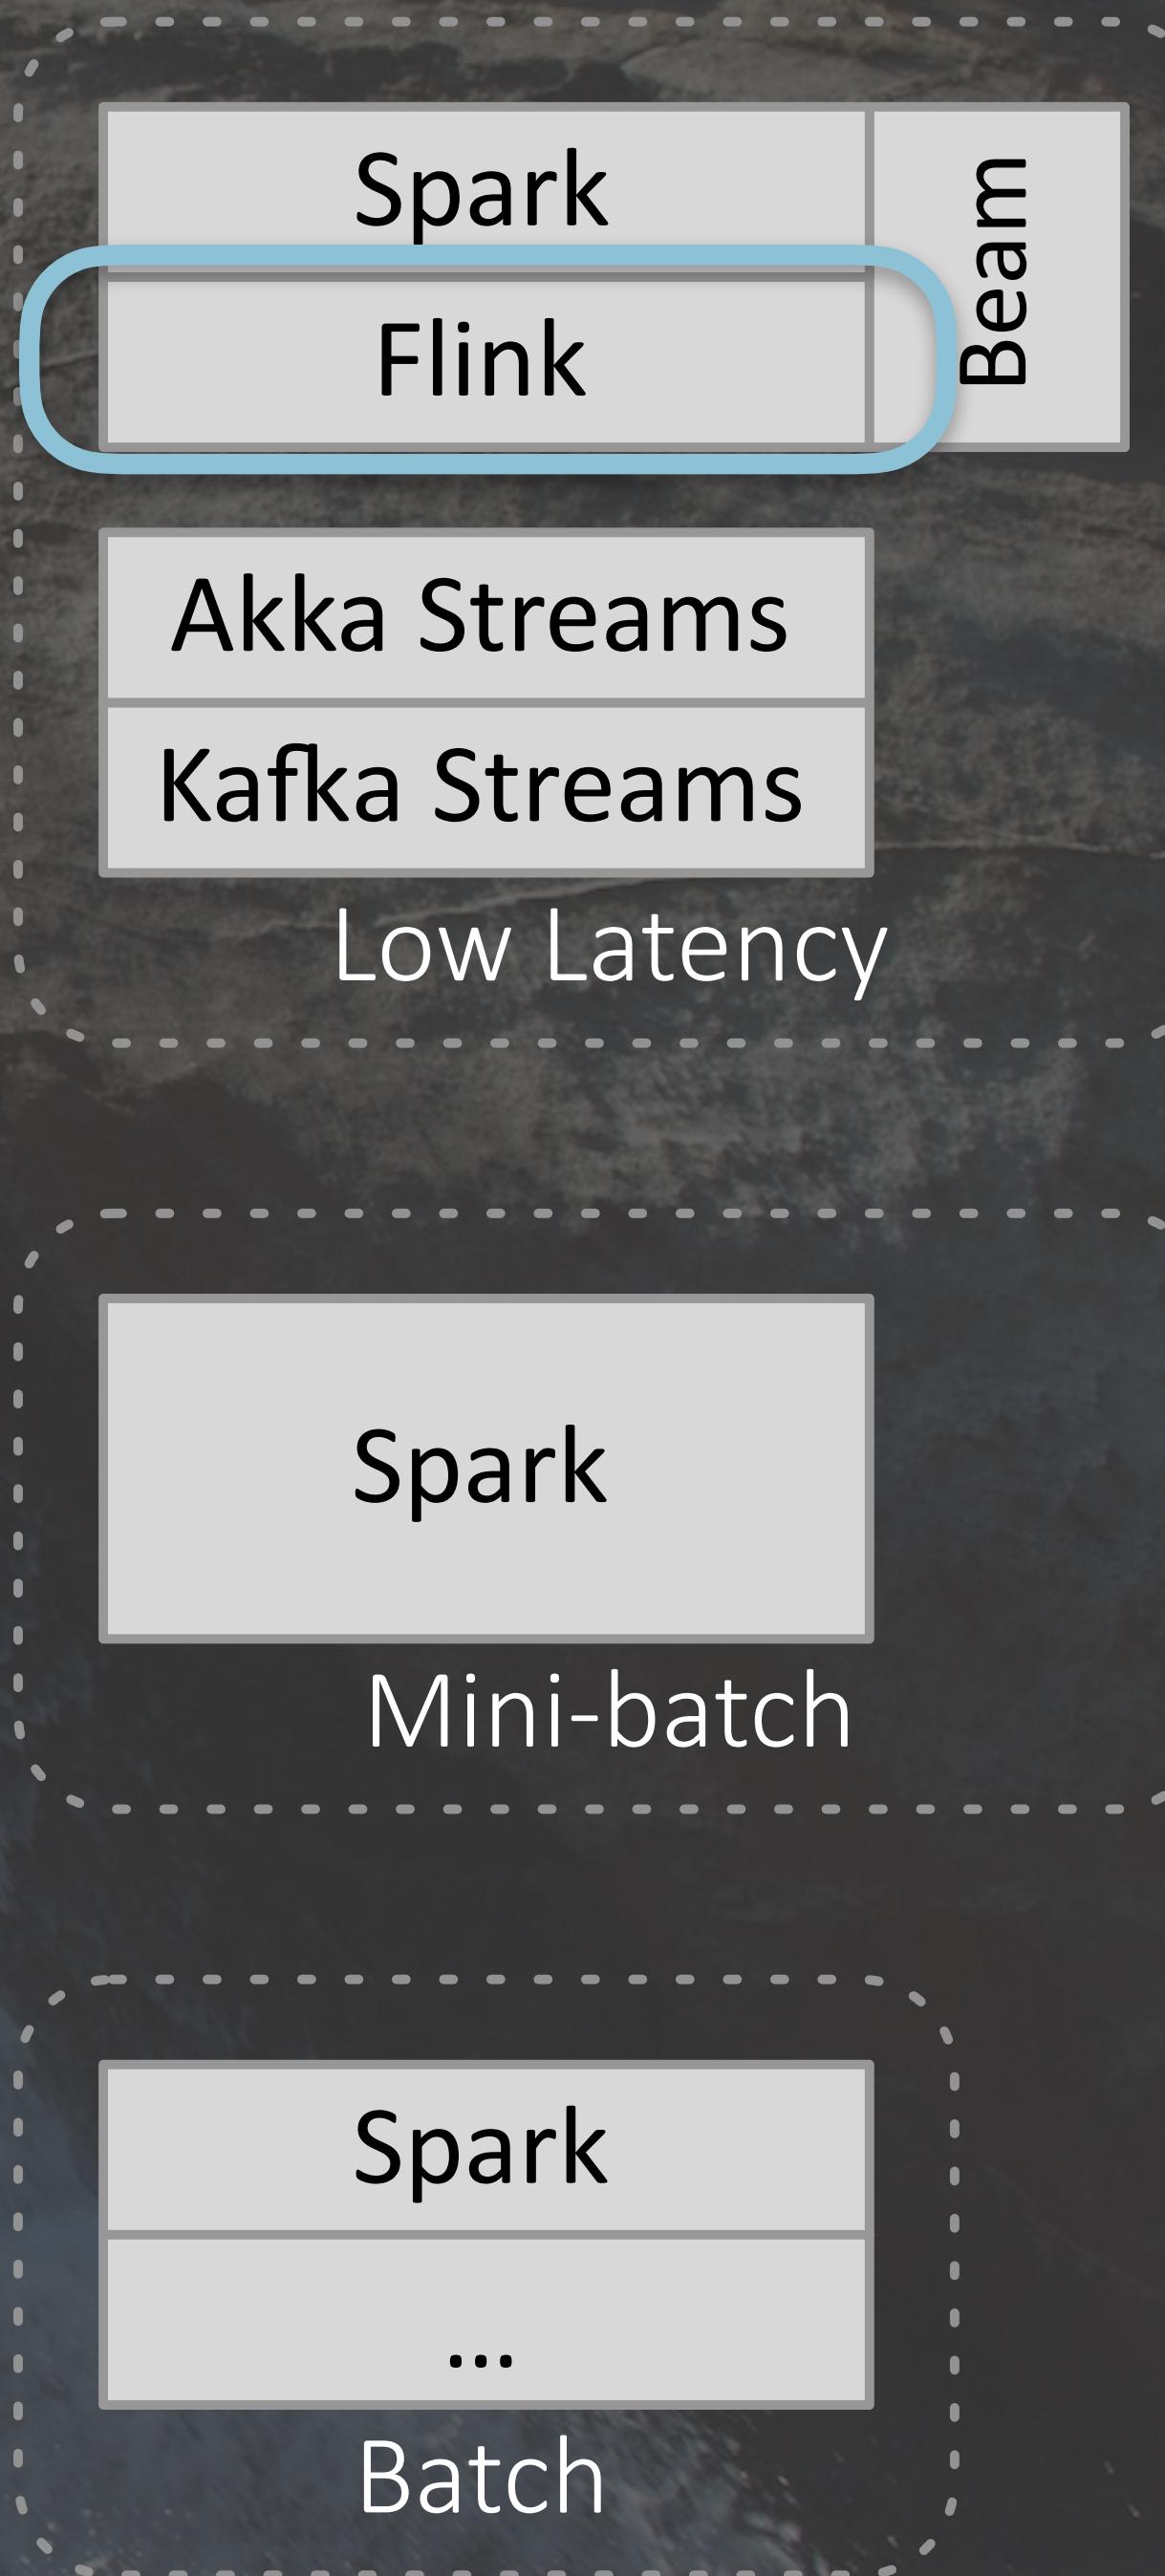
Spark

- Spark Batch
- Same Dataset and RDD features as streaming.
- Massive scalability
- Excellent performance

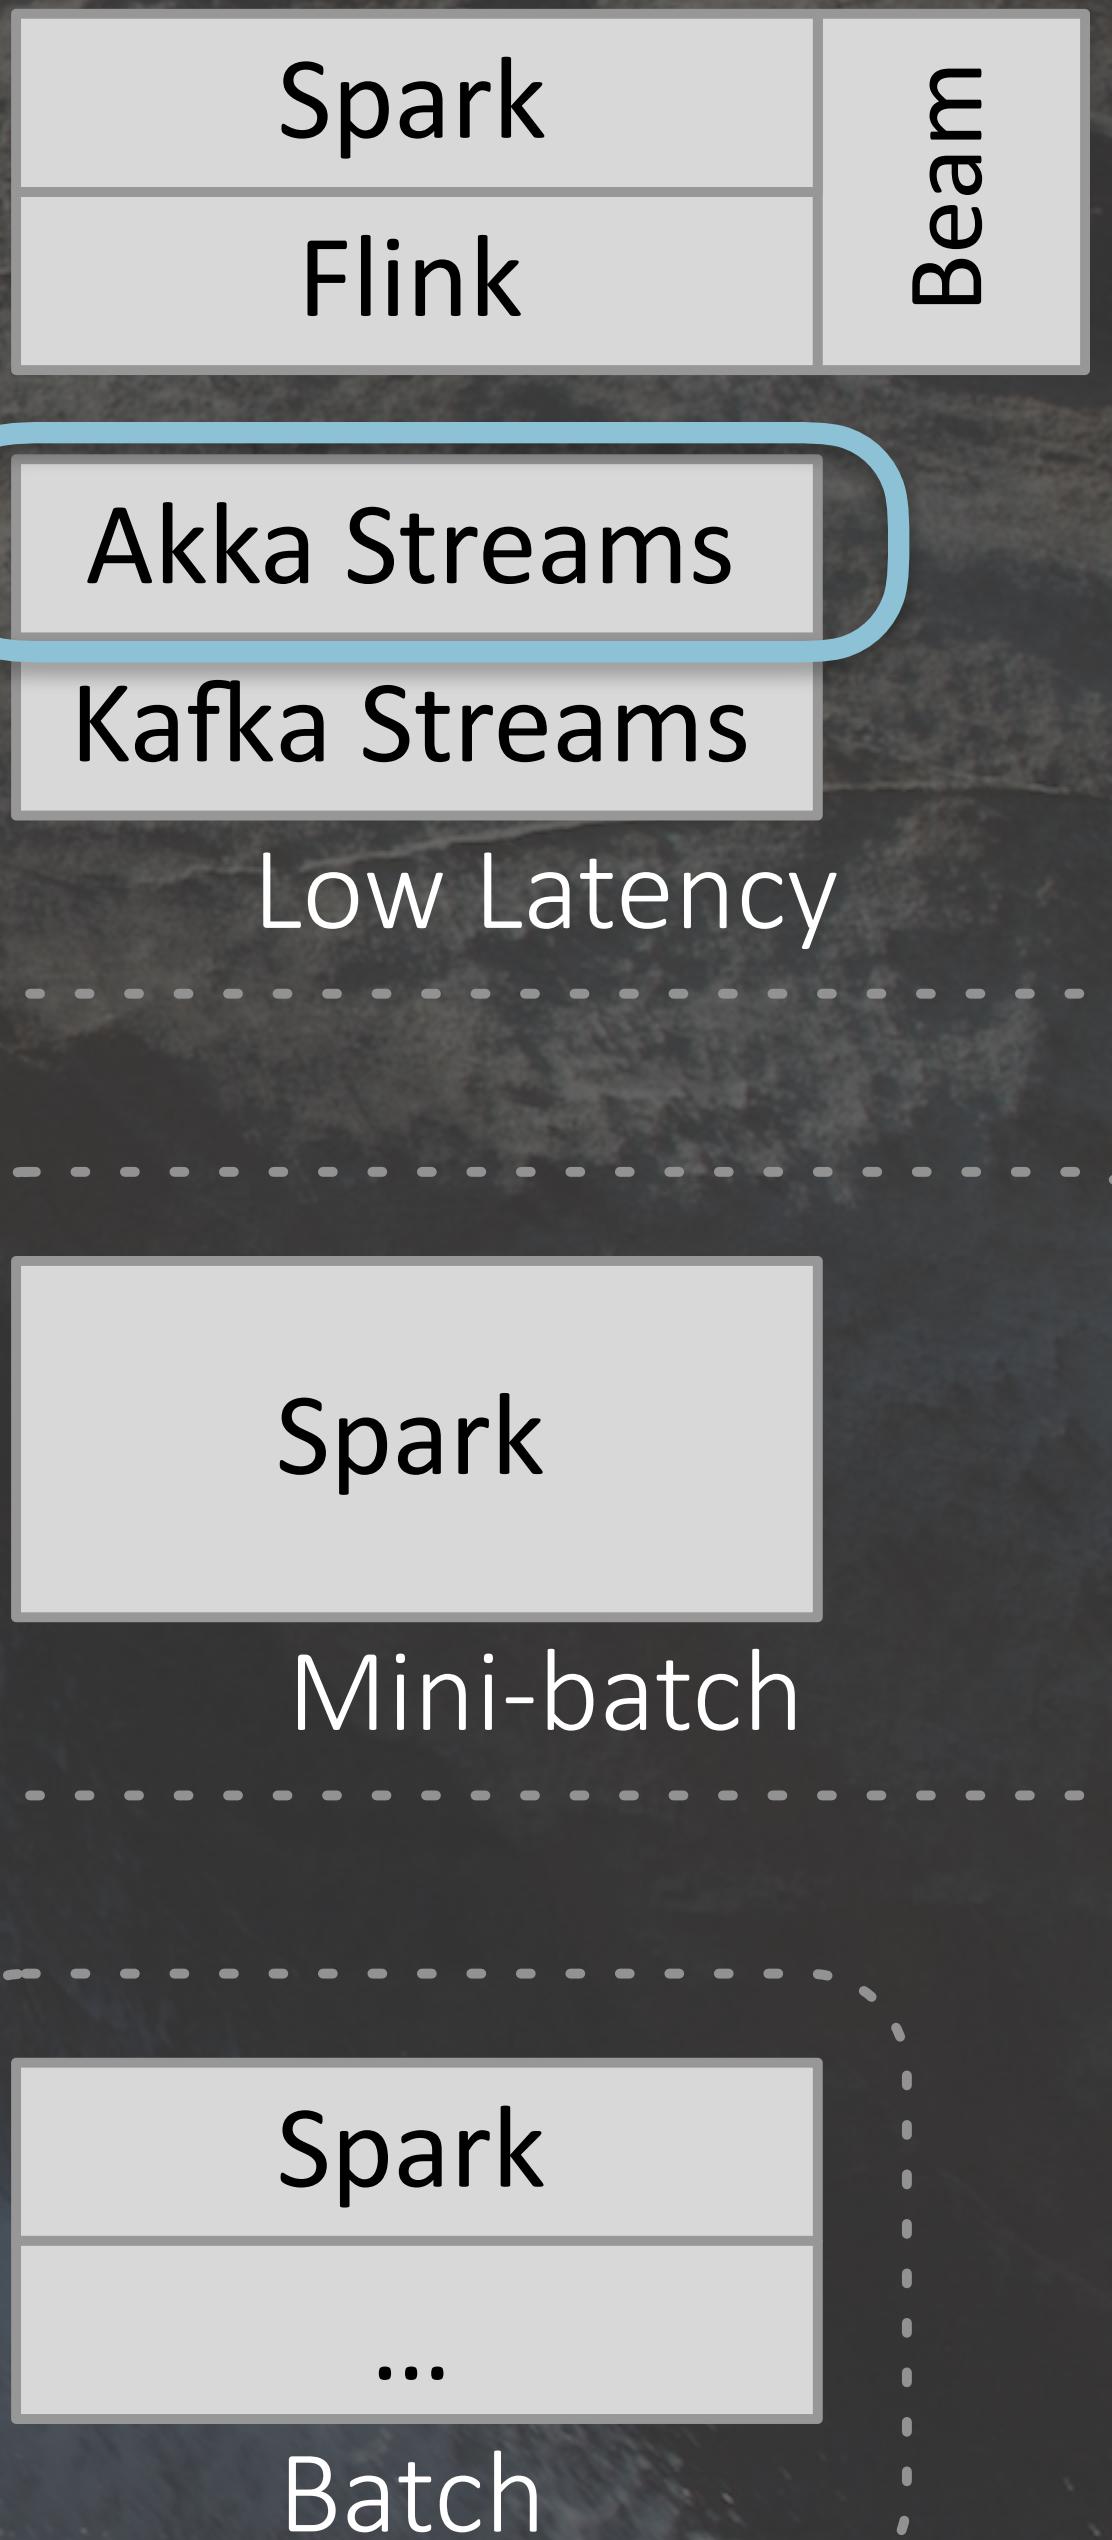


Spark 

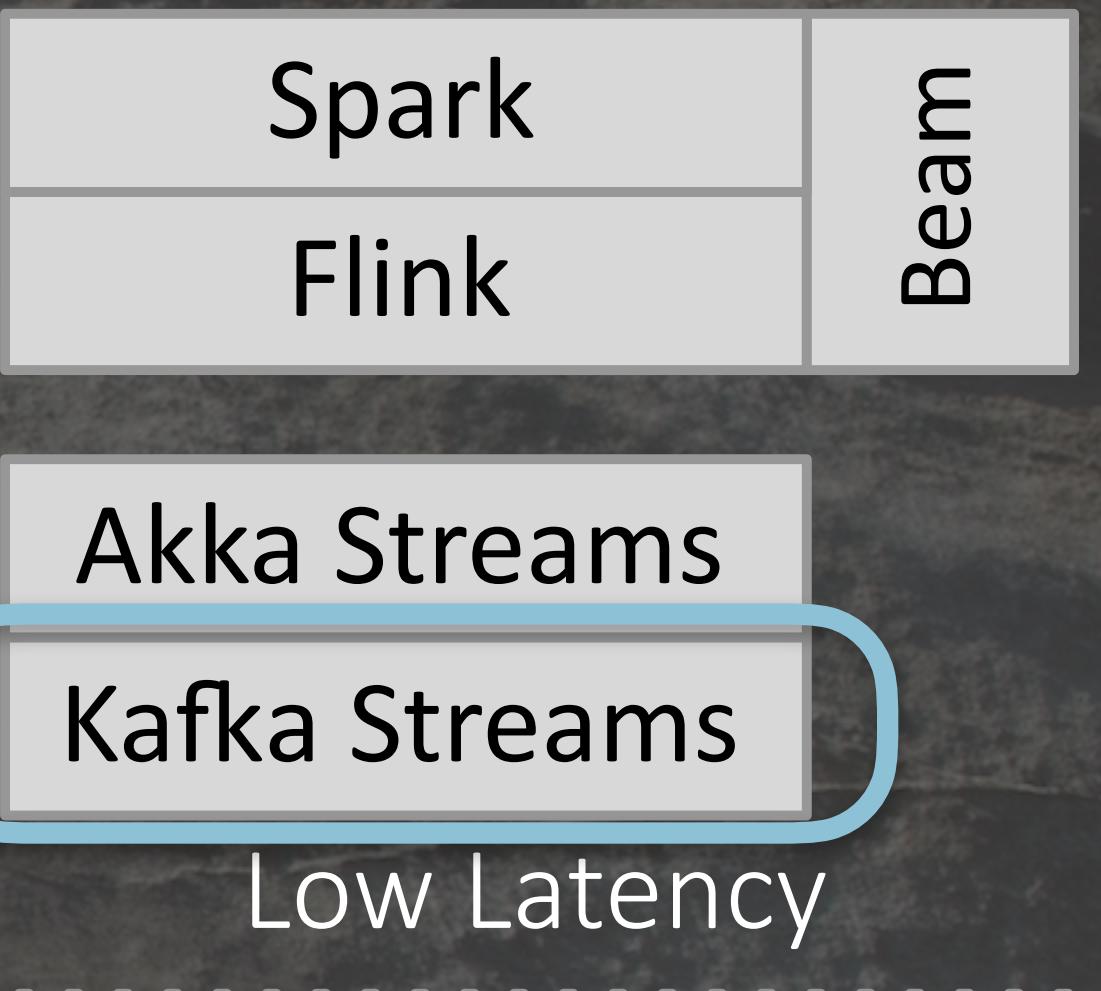
- Apache Flink
- High volume, low latency
- Sophisticated streaming (Beam) semantics
- SQL, evolving ML support



- Akka Streams
- Low latency
- Complex Event Processing
- Efficient, per event
- Mid-volume pipelines



- Kafka Streams
- Low overhead Kafka topic processing
- Ideal for ETL and aggregations



Spark

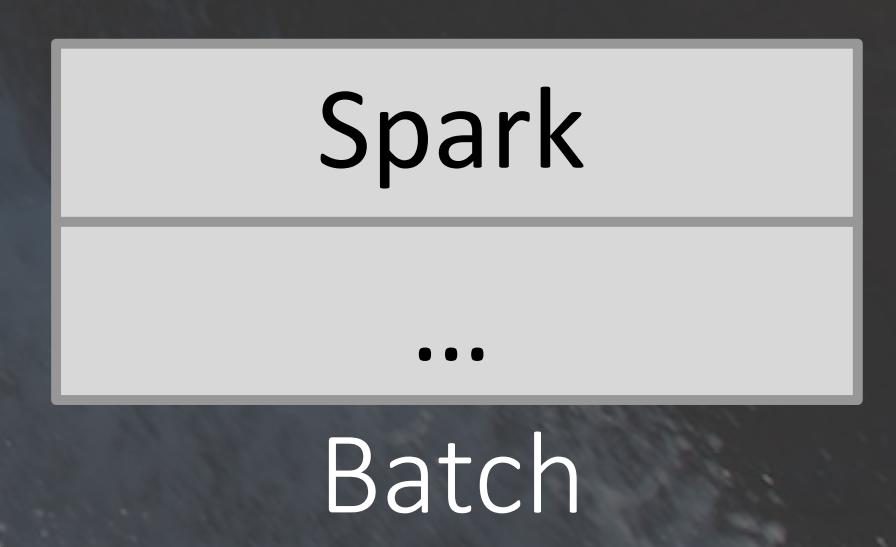
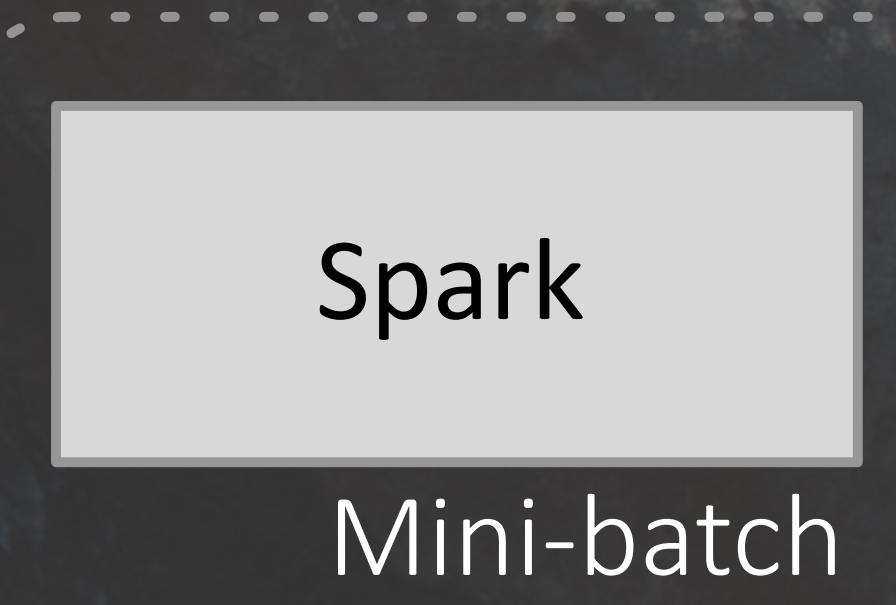
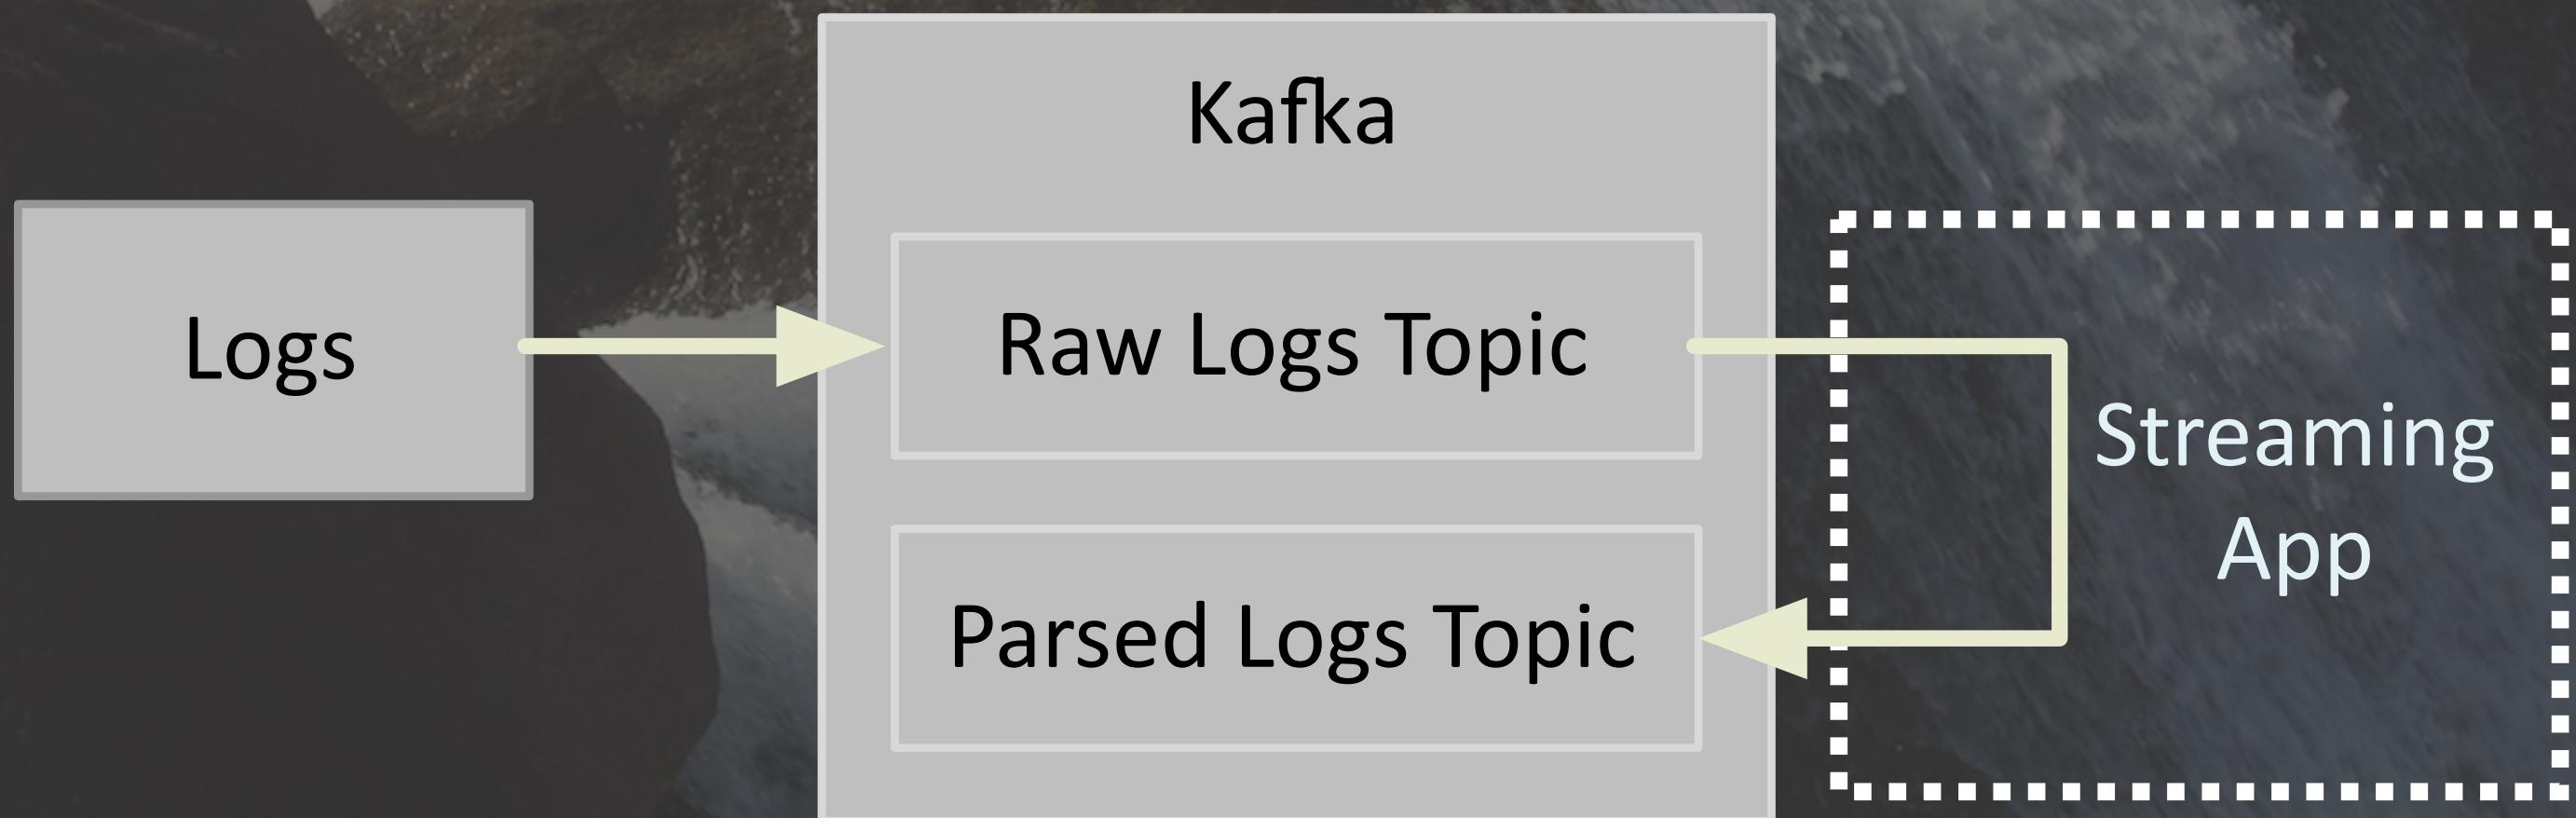
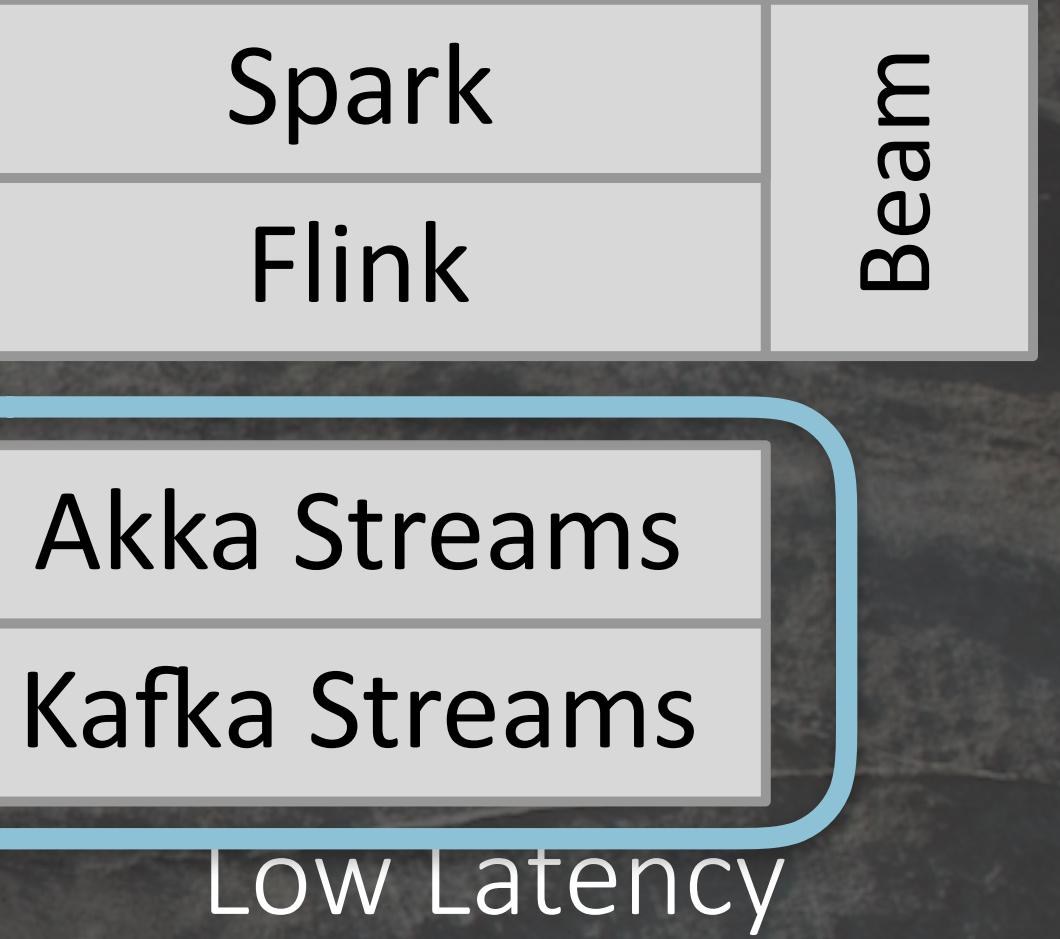
Mini-batch

Spark

...

Batch

- Akka and Kafka Streams
- “Exactly once” with transactions



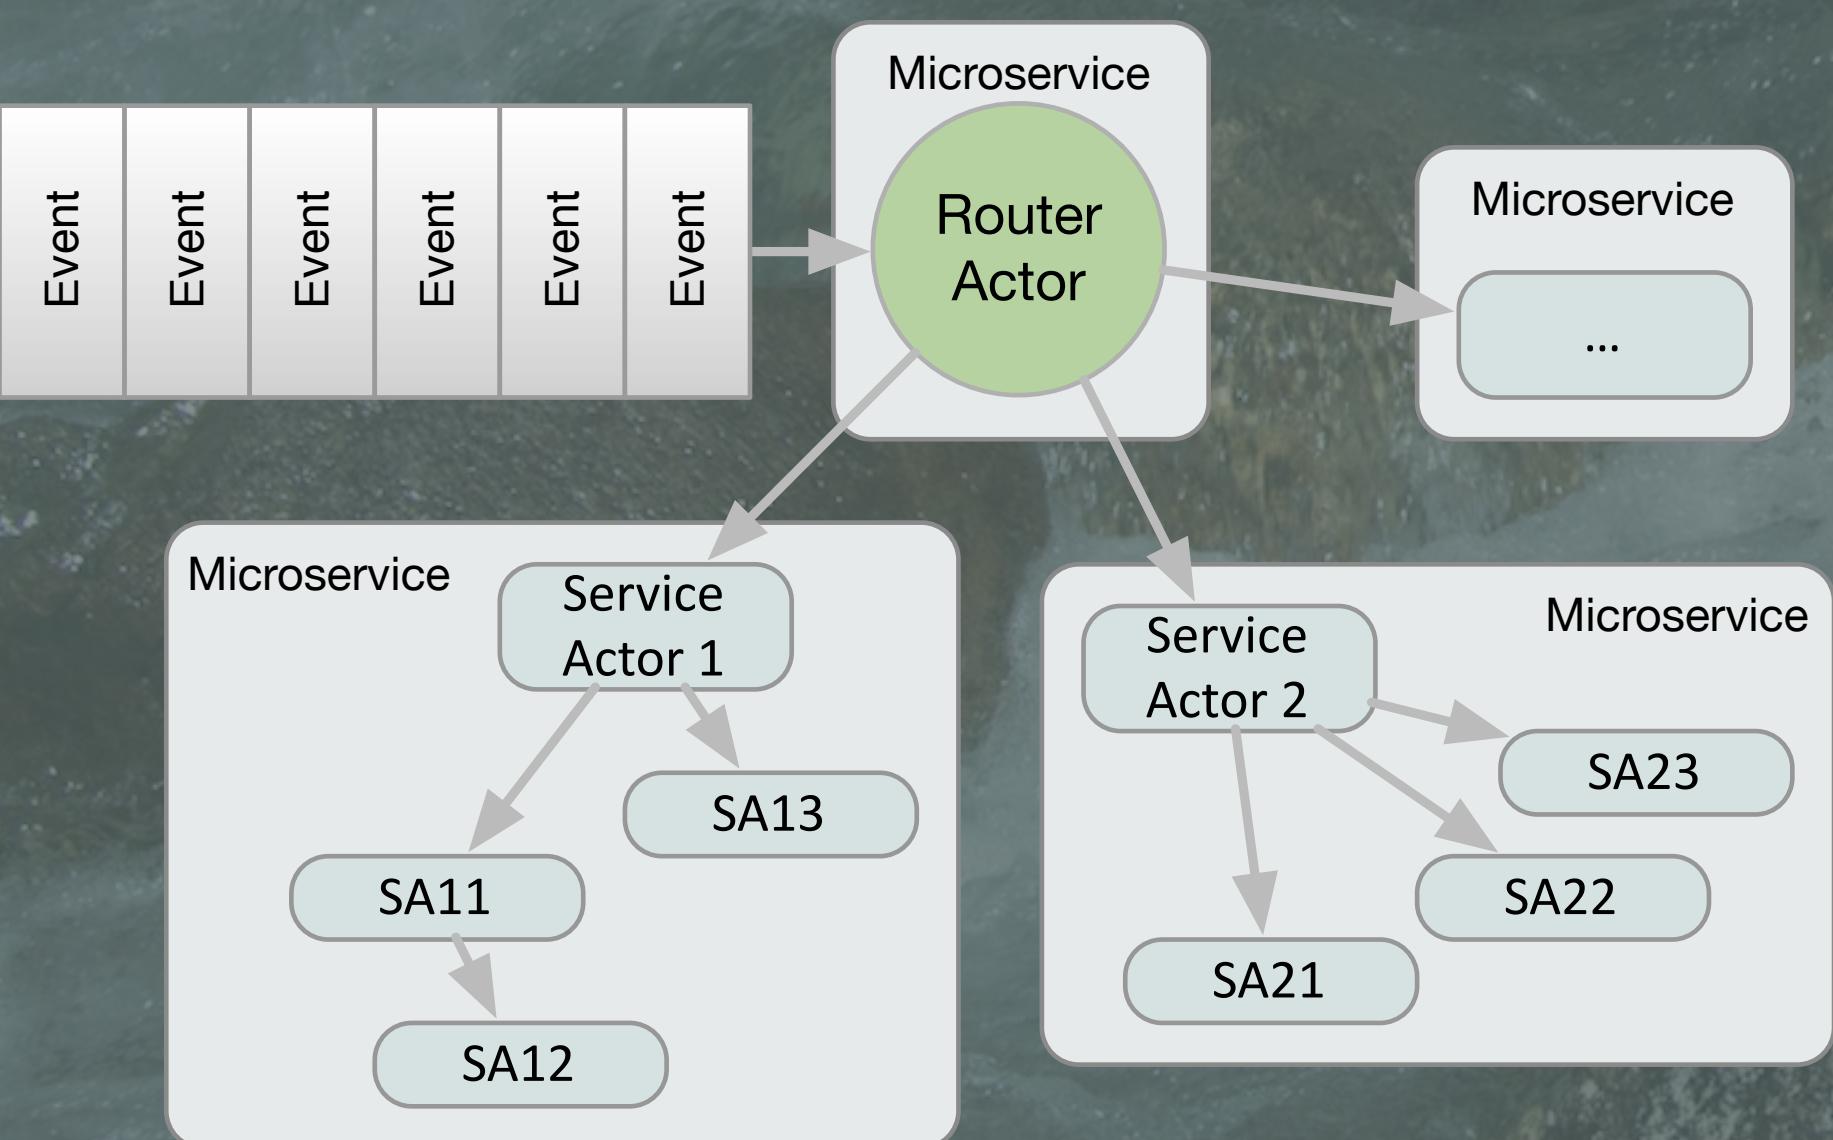
• Process data individually or in bulk?



Each grew out of one end of this spectrum...

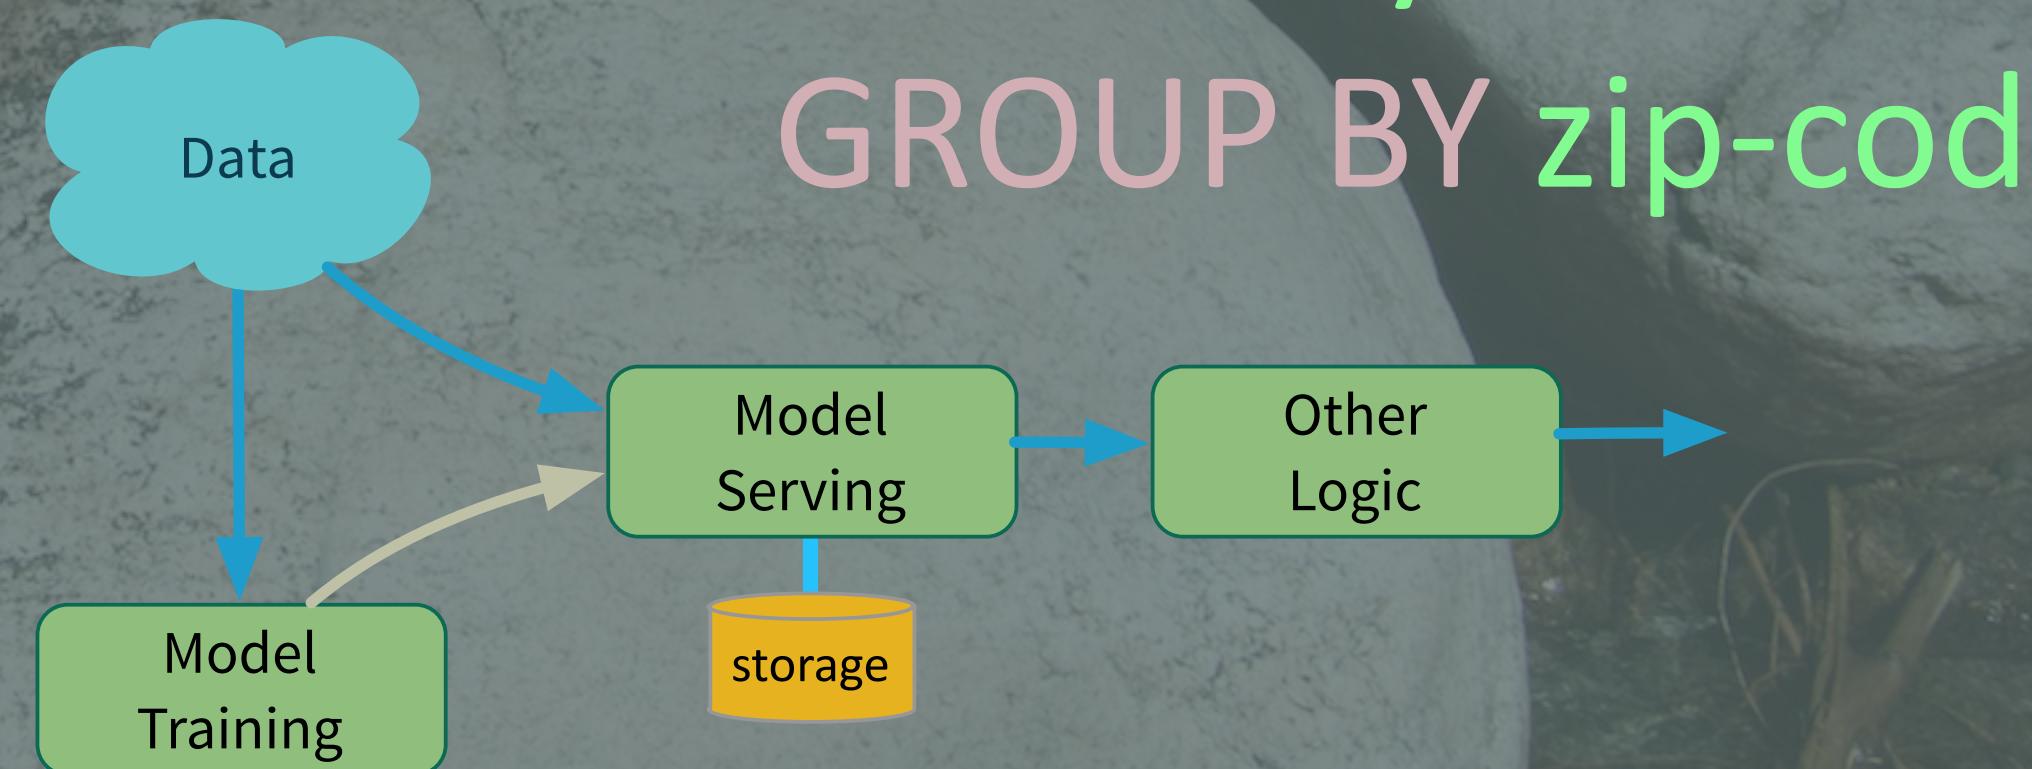


Event-driven μ -services



“Record-centric” μ -services

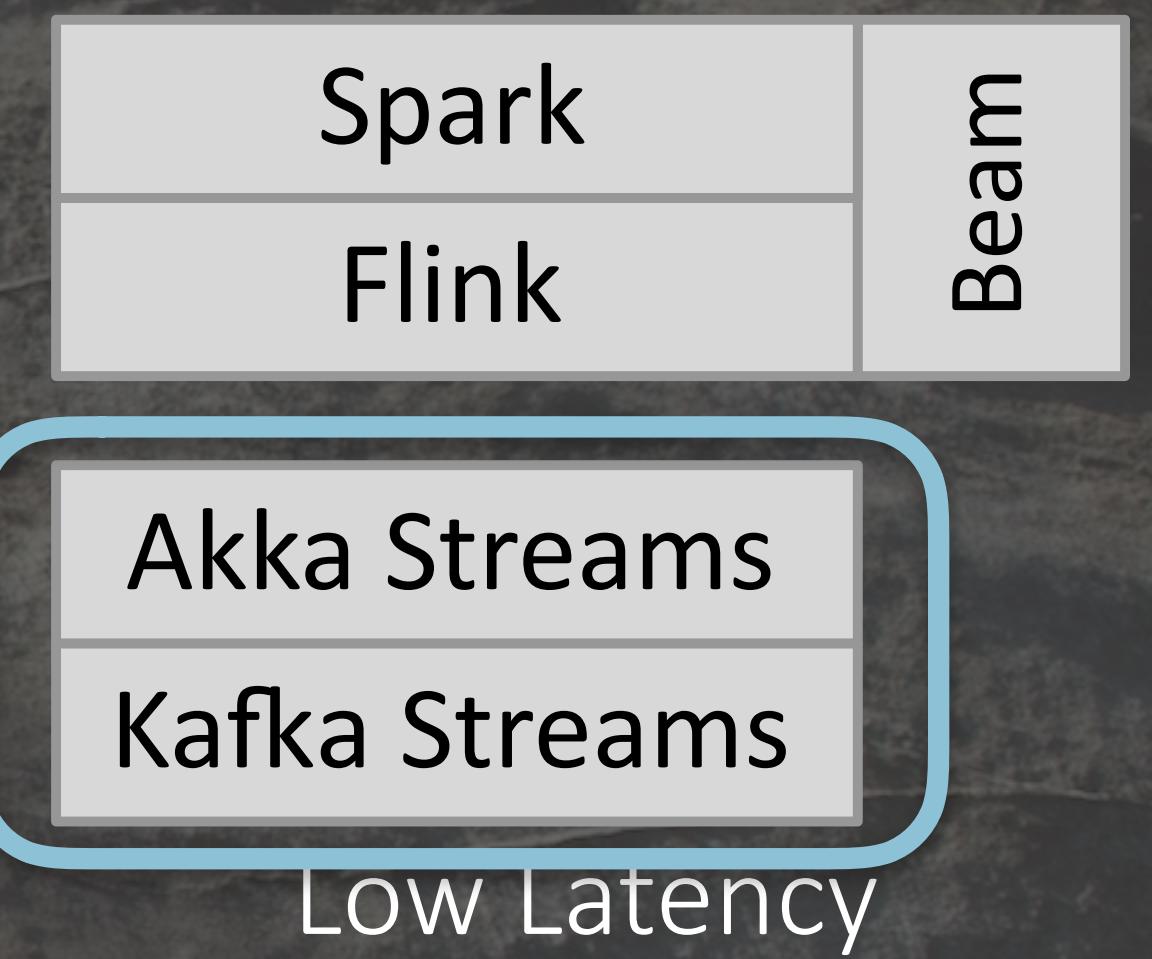
`SELECT COUNT(*)
FROM my-iot-data
GROUP BY zip-code`



Events

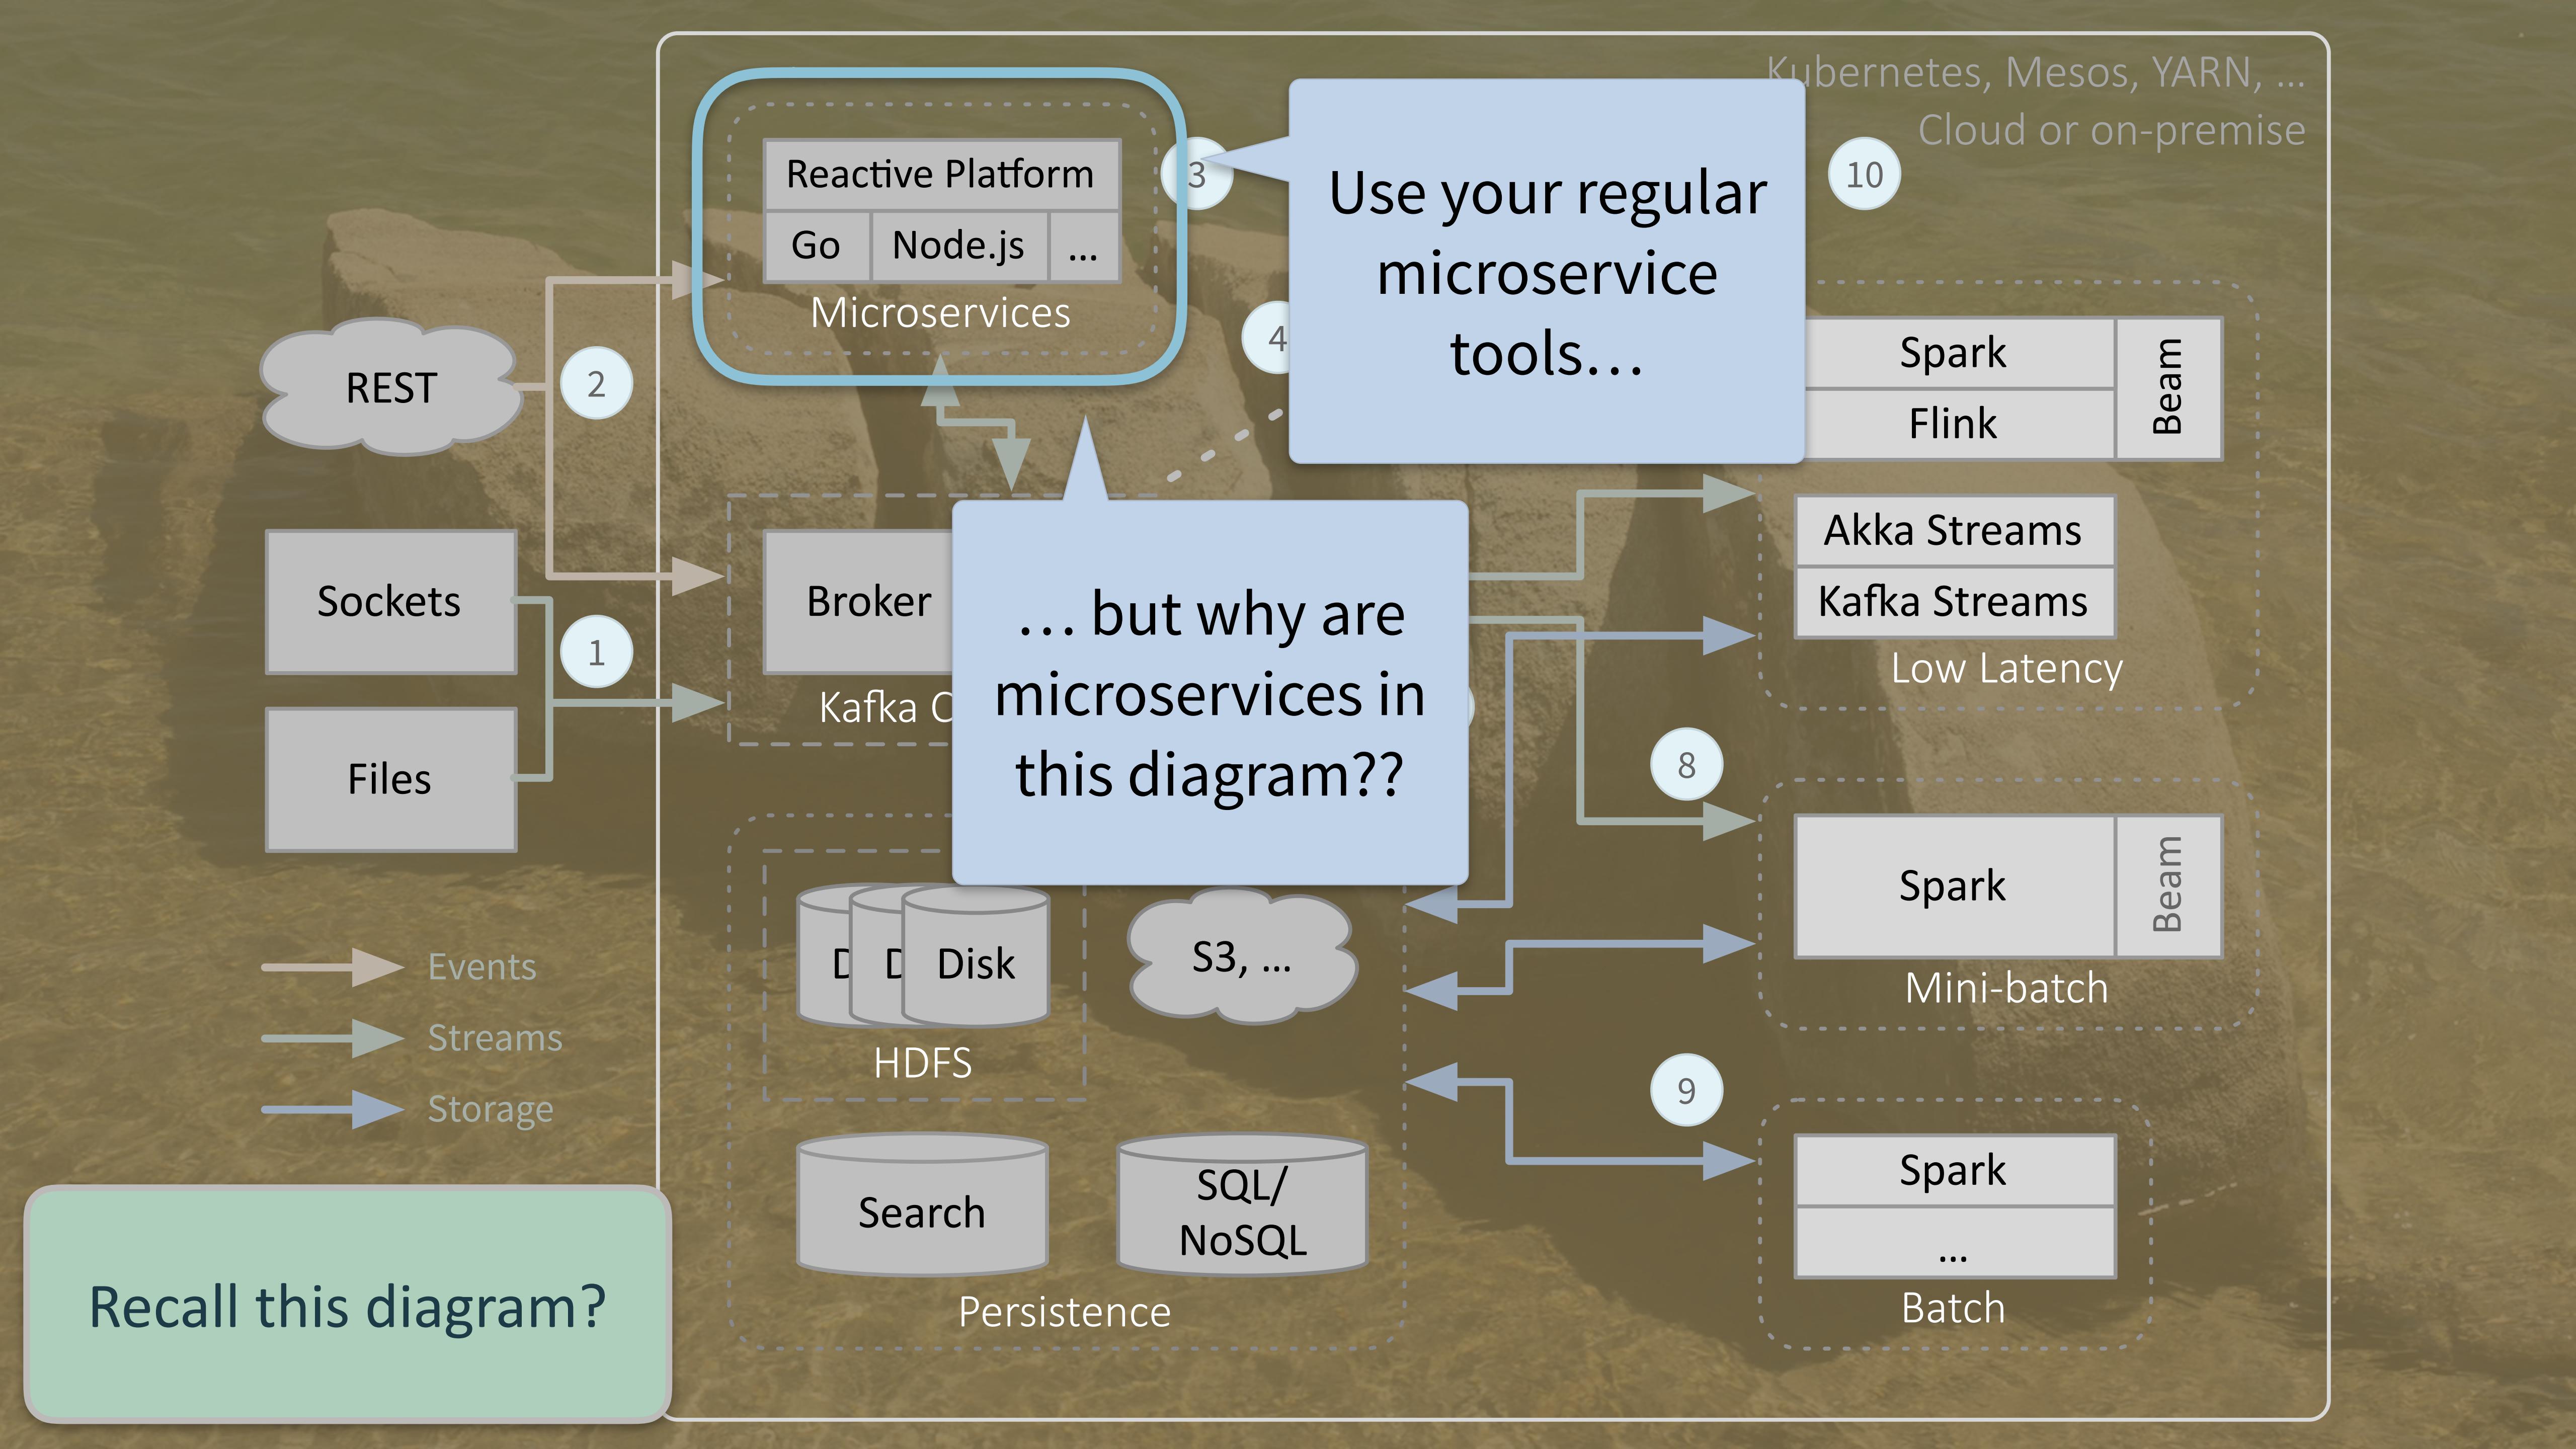
Records

- Akka Streams vs. Kafka Streams talk
- Also at polyglotprogramming.com/talks/

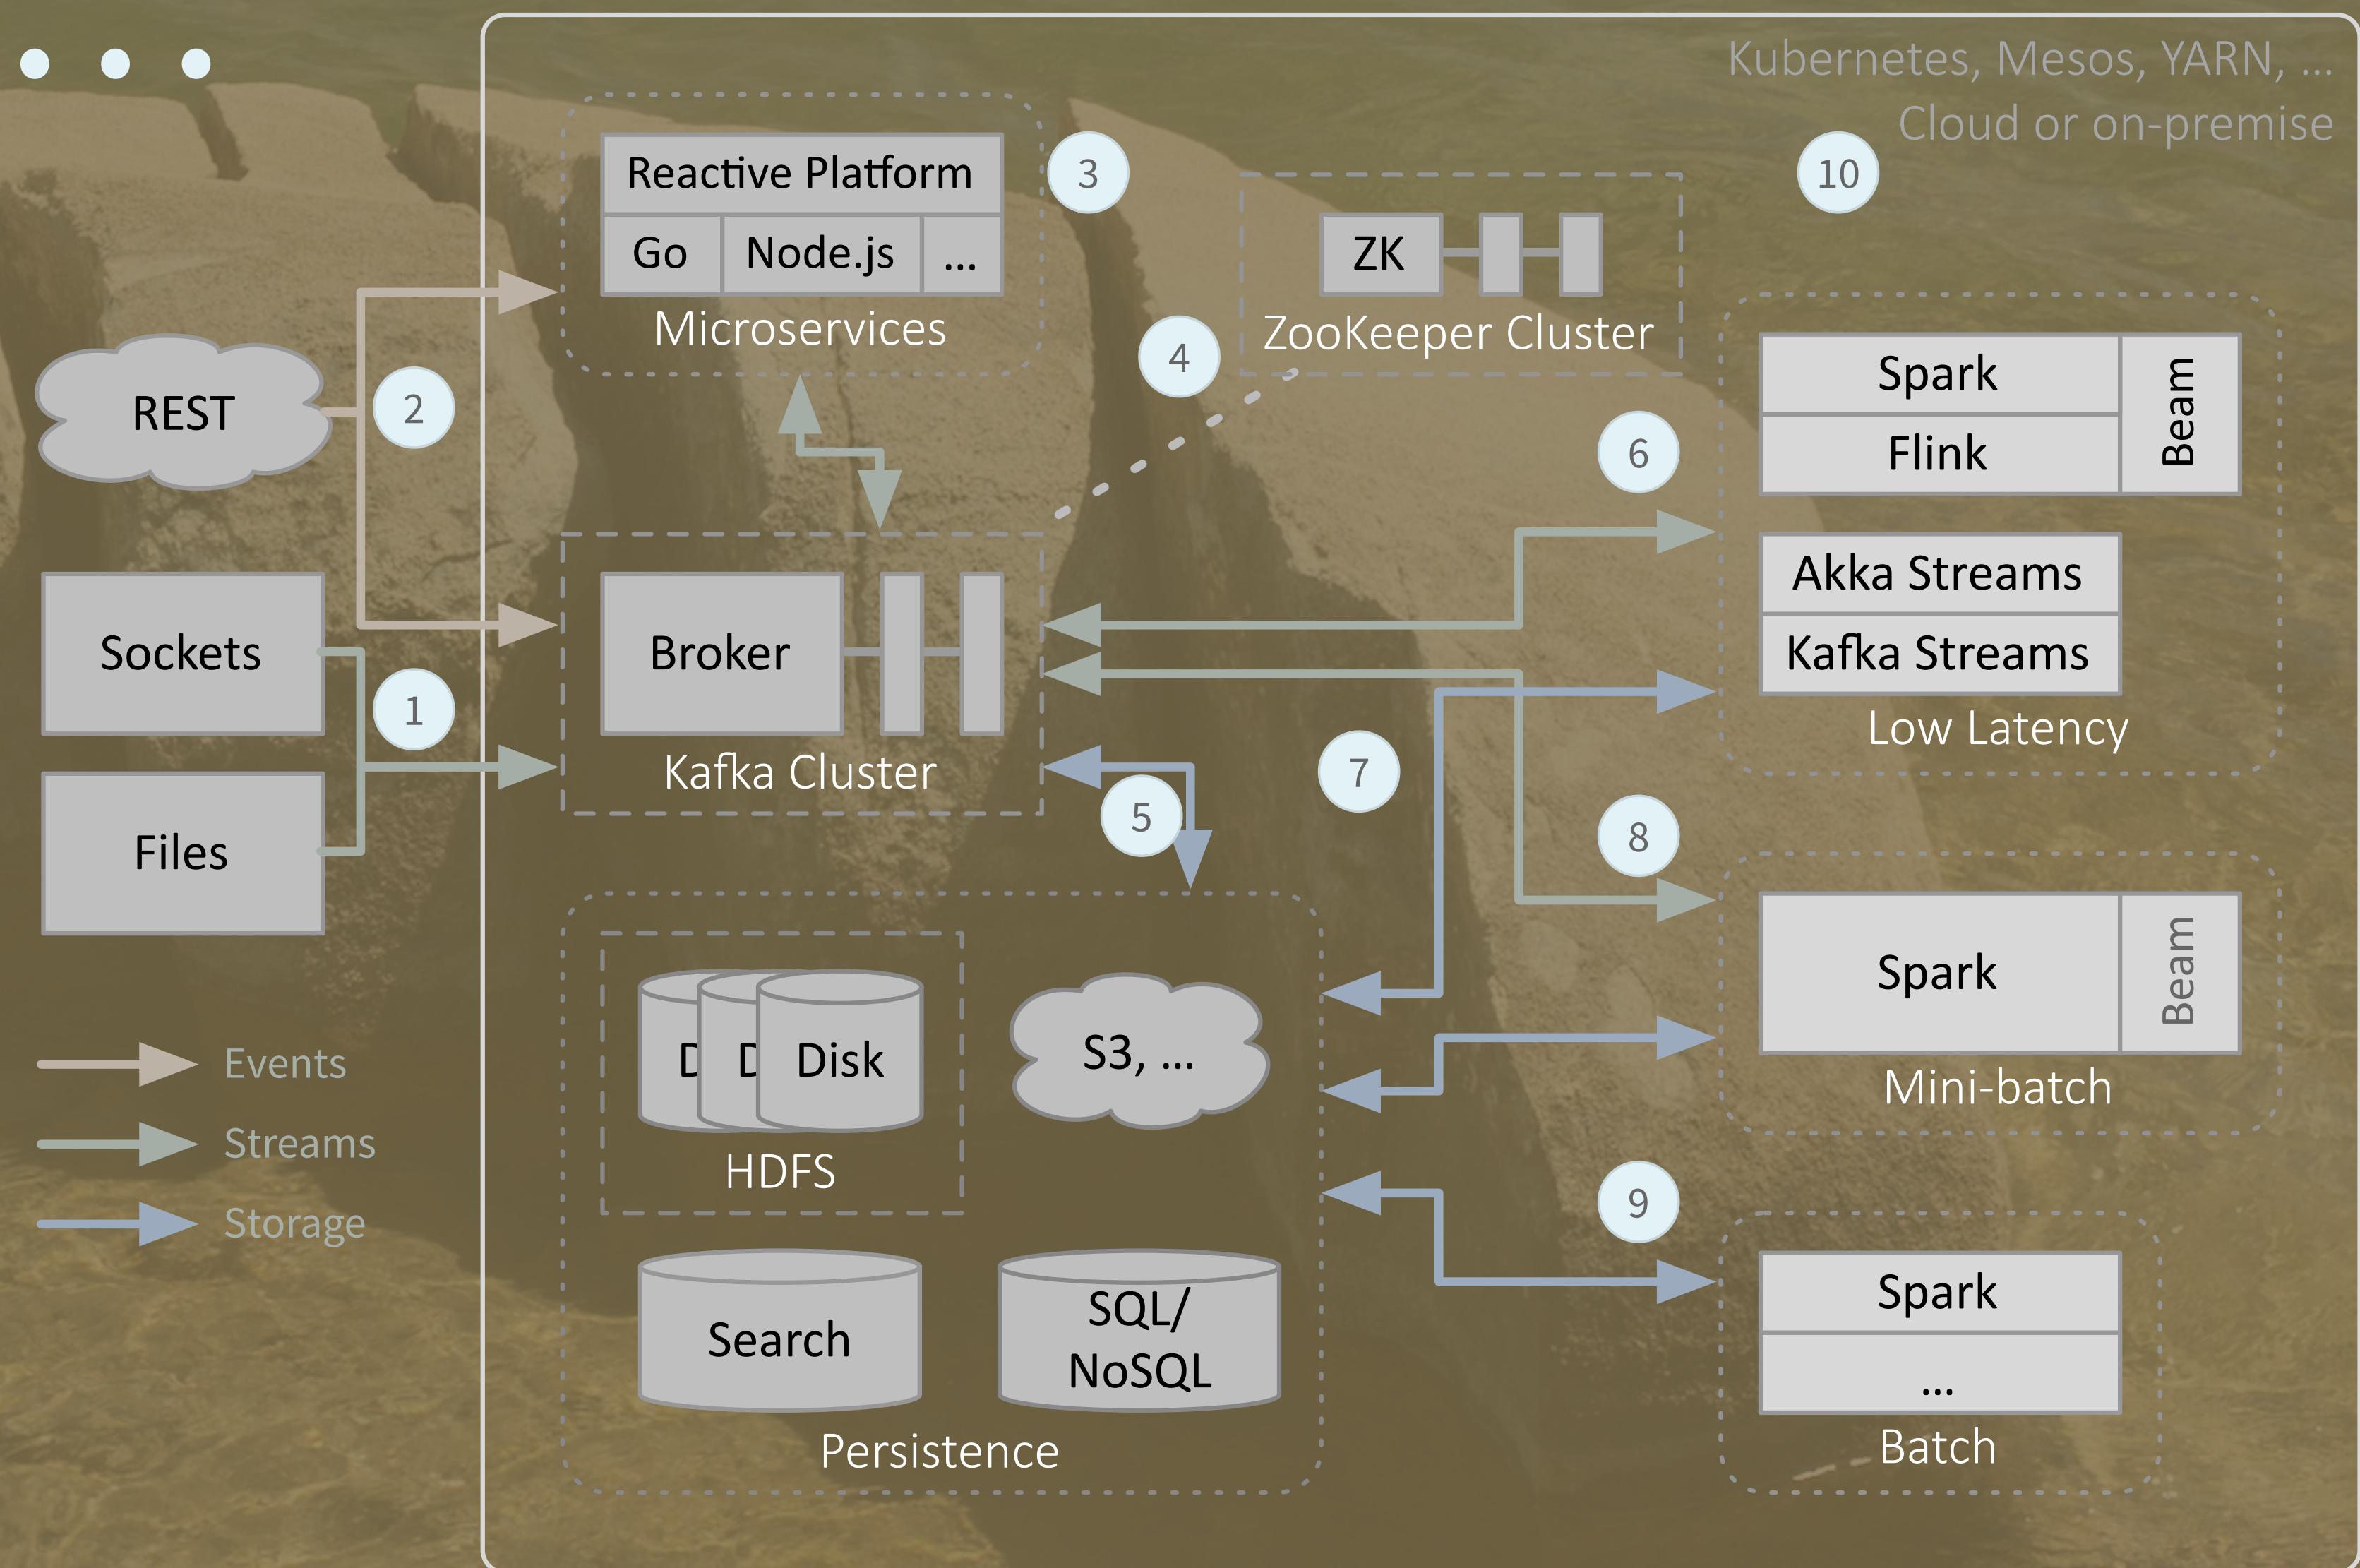


A large, light-colored rock formation, possibly granite, is partially submerged in clear, shallow water. The rocks are angular and have sharp edges. In the background, there are greenish-blue waves, suggesting a lake or coastal area. The lighting is bright, casting shadows on the rocks.

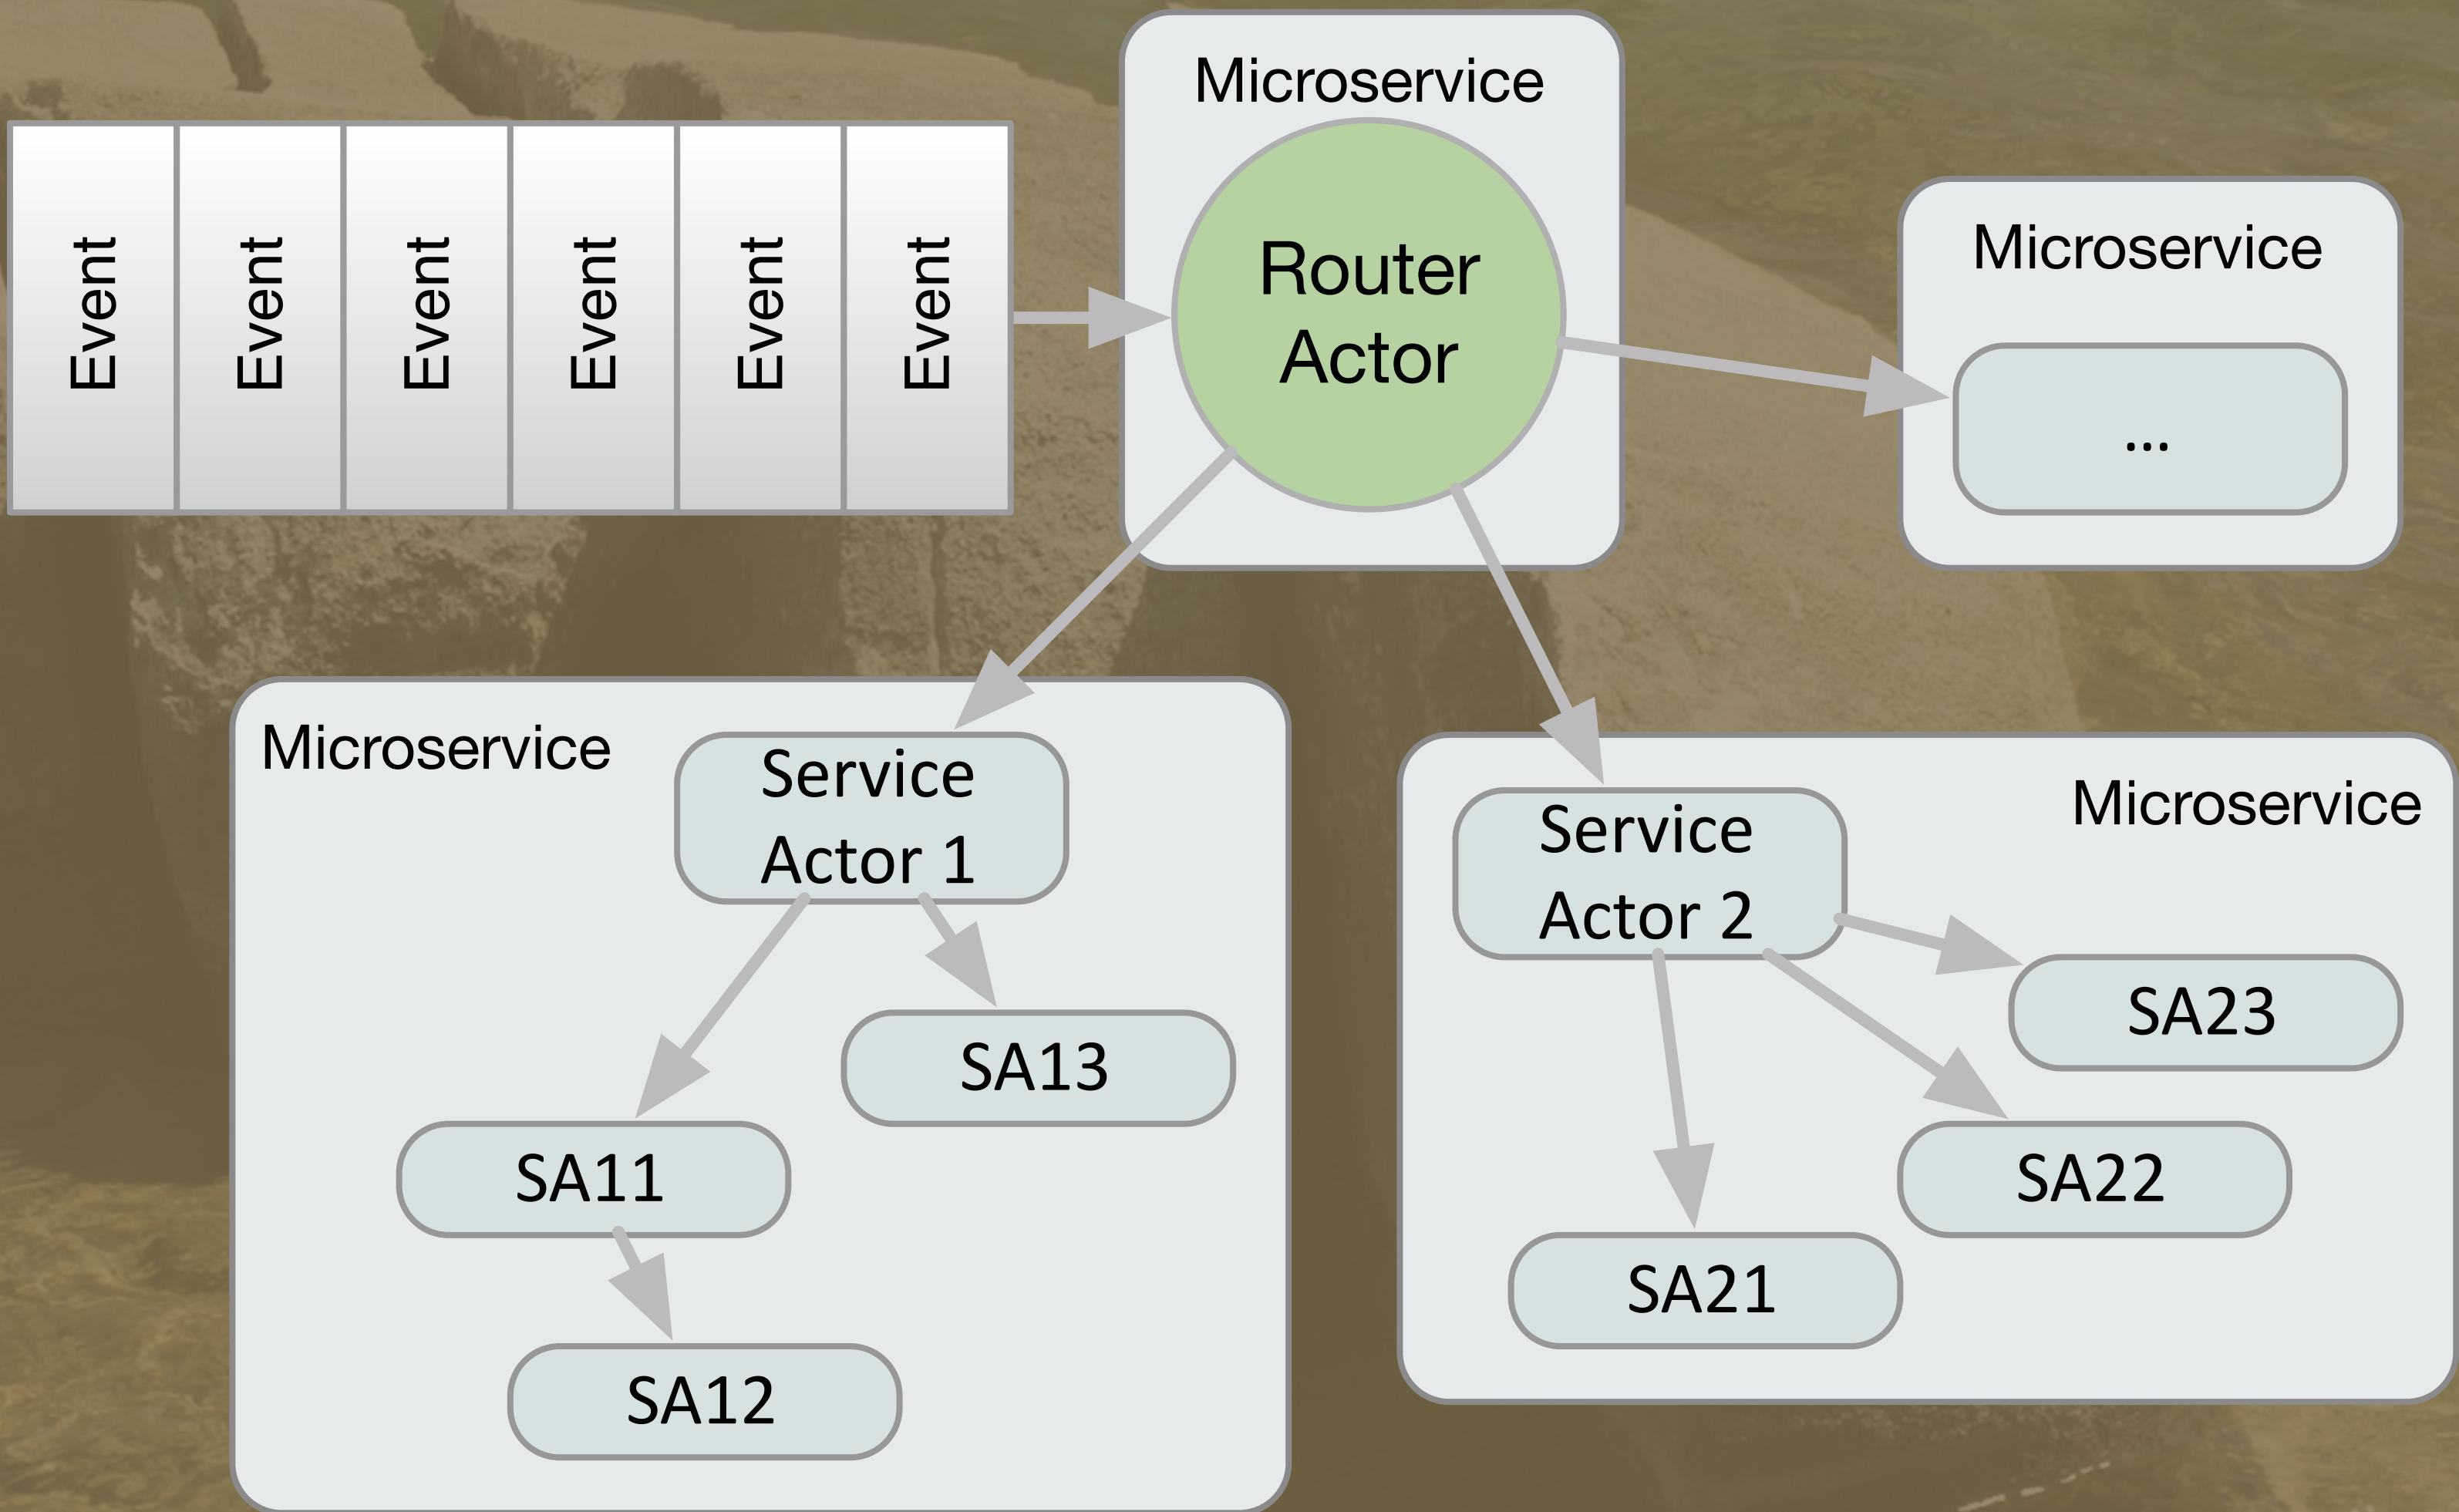
Microservices and Fast Data



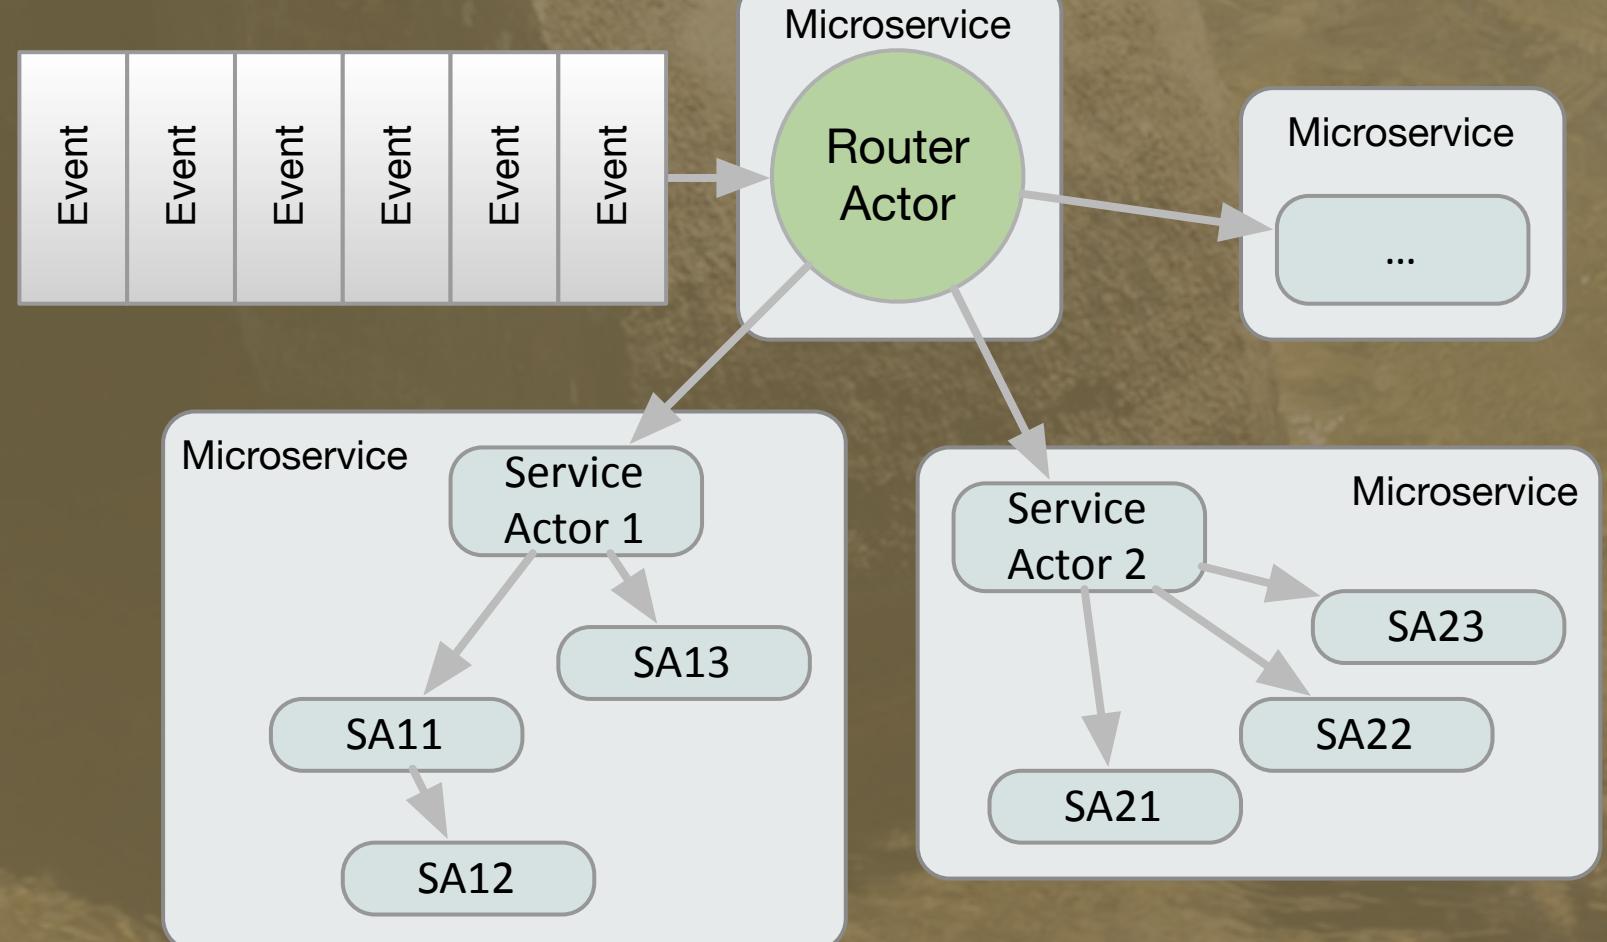
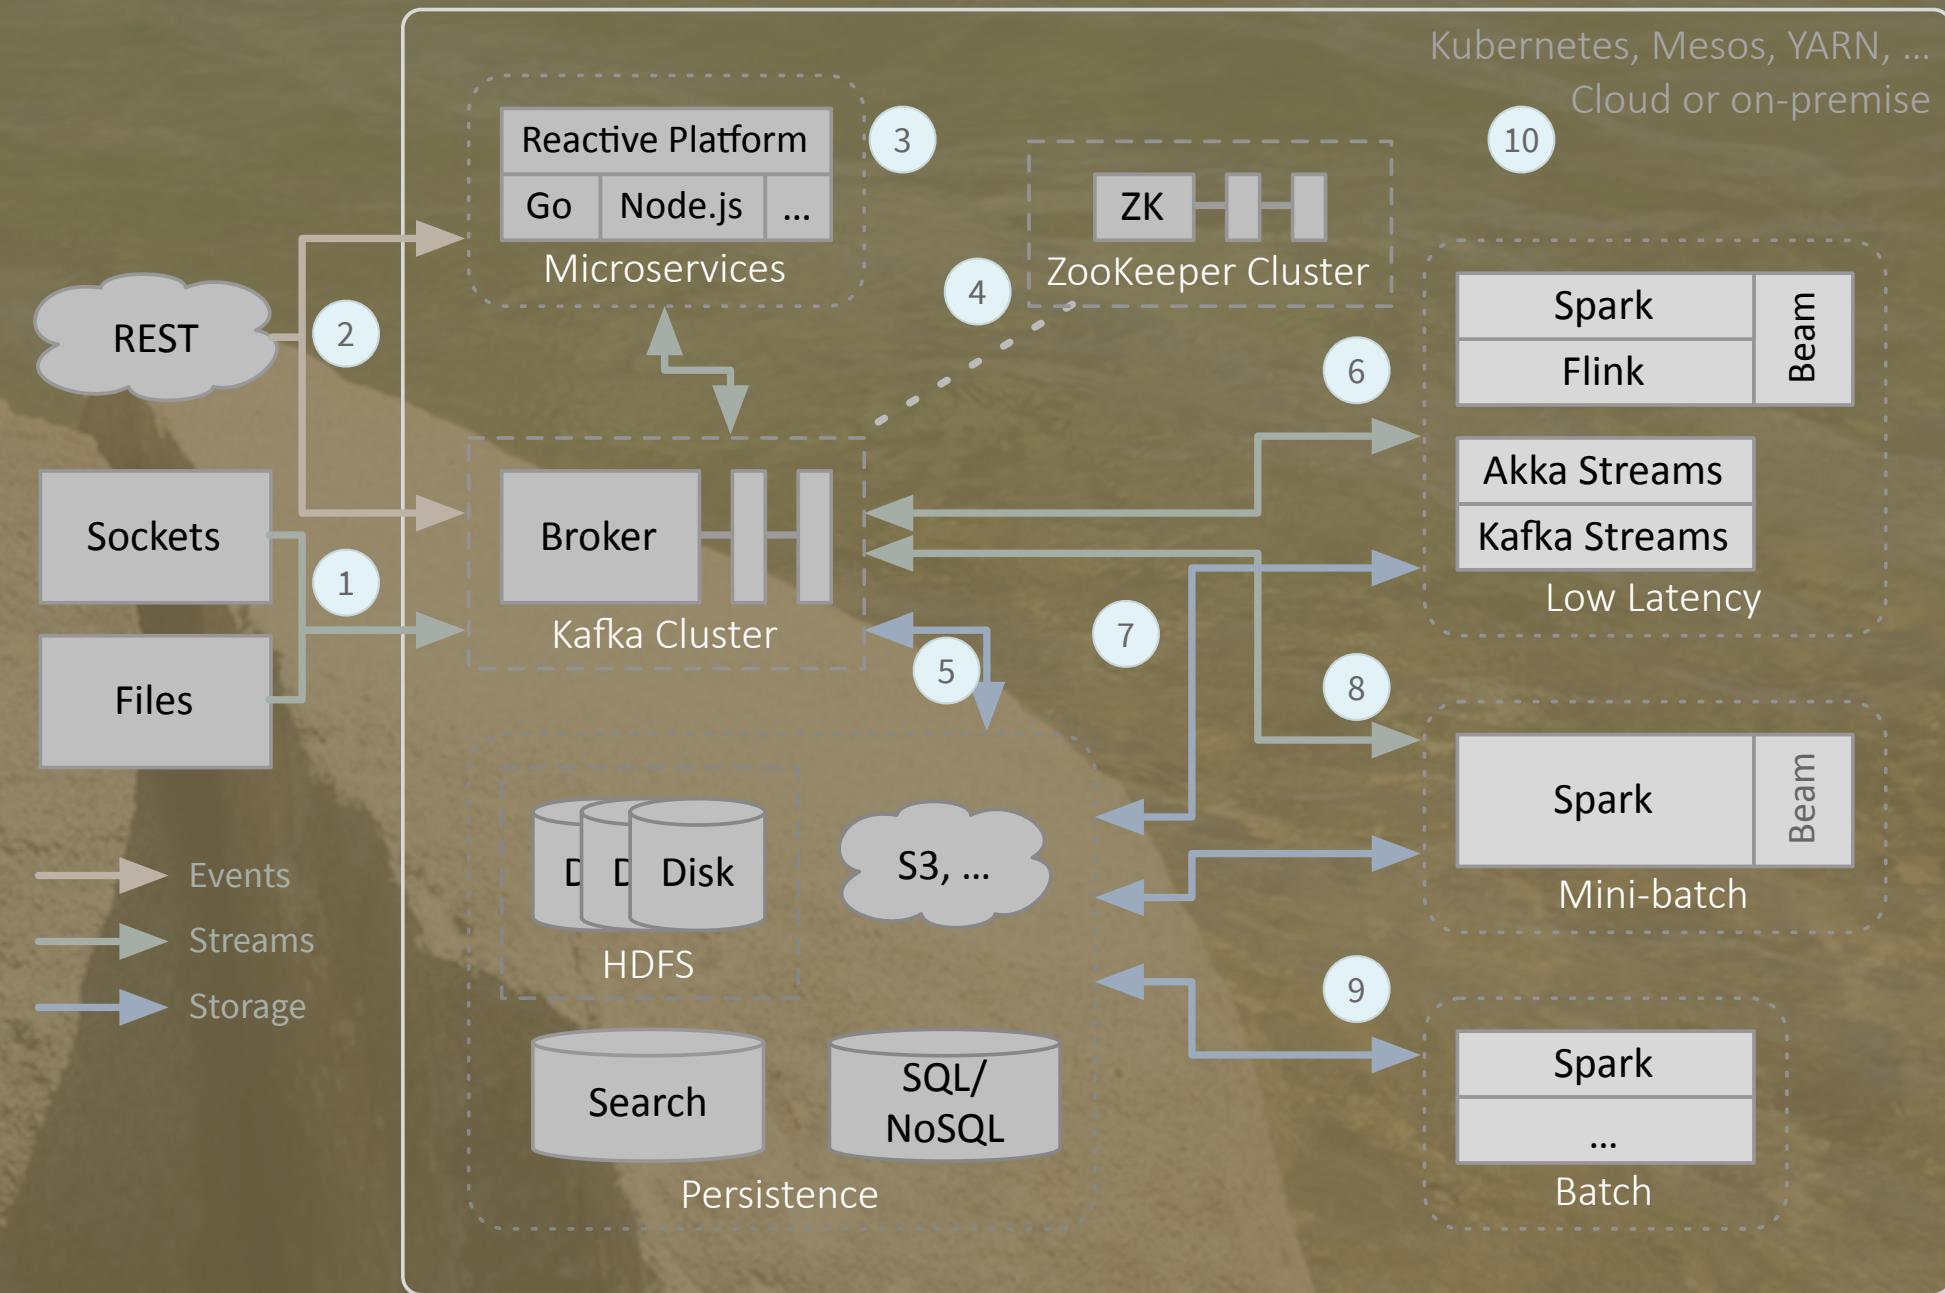
How is this...



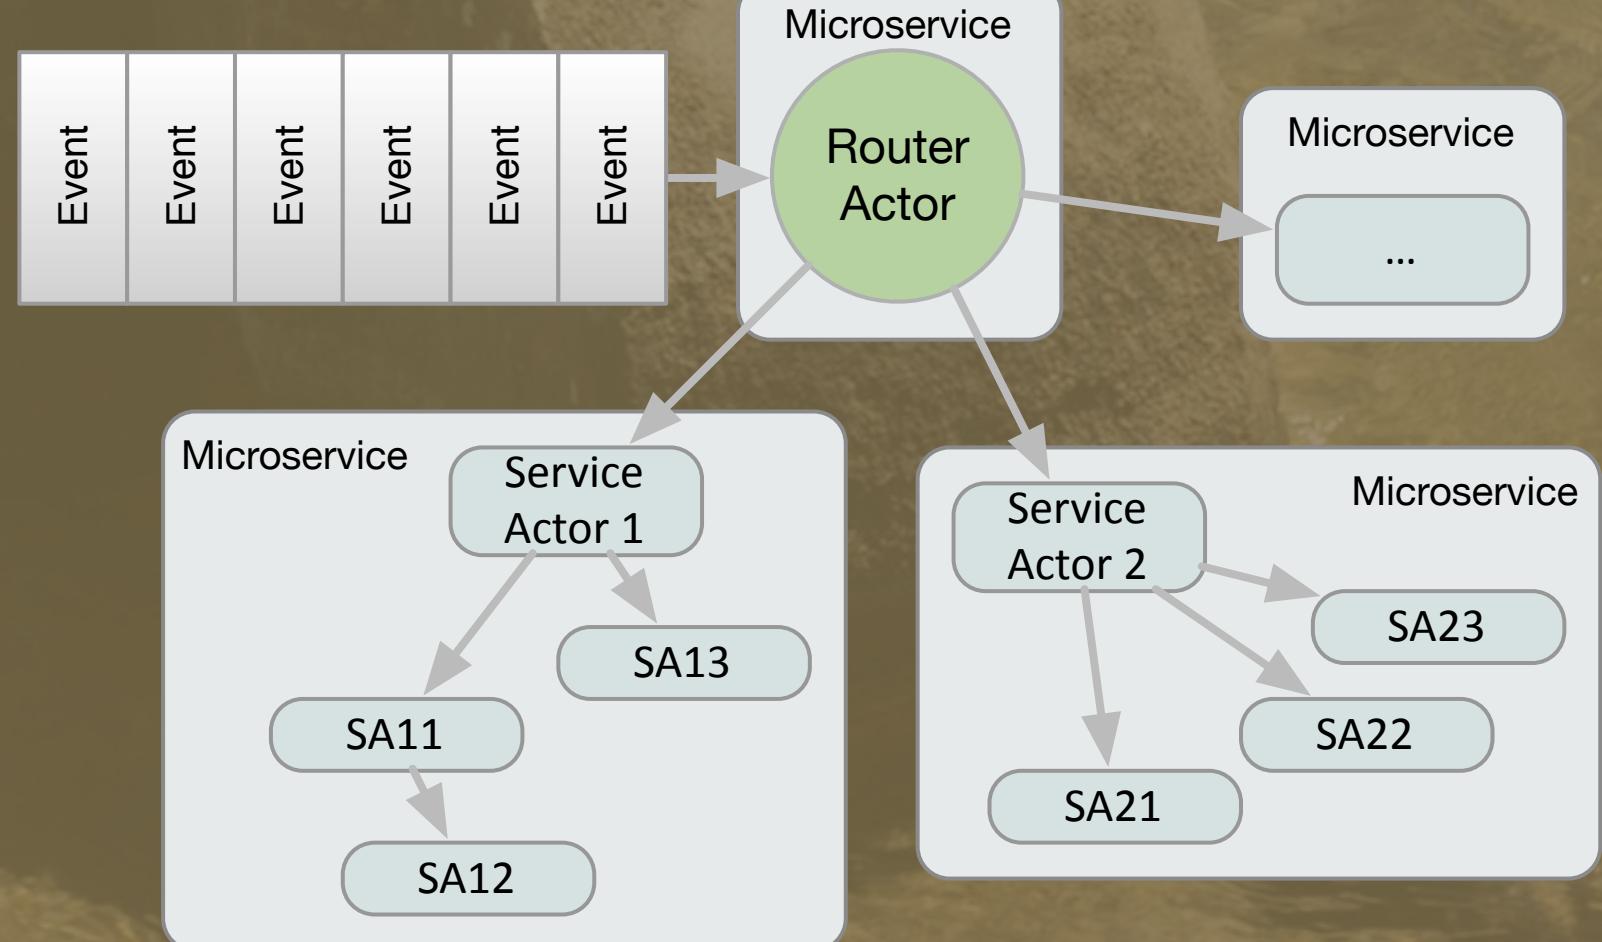
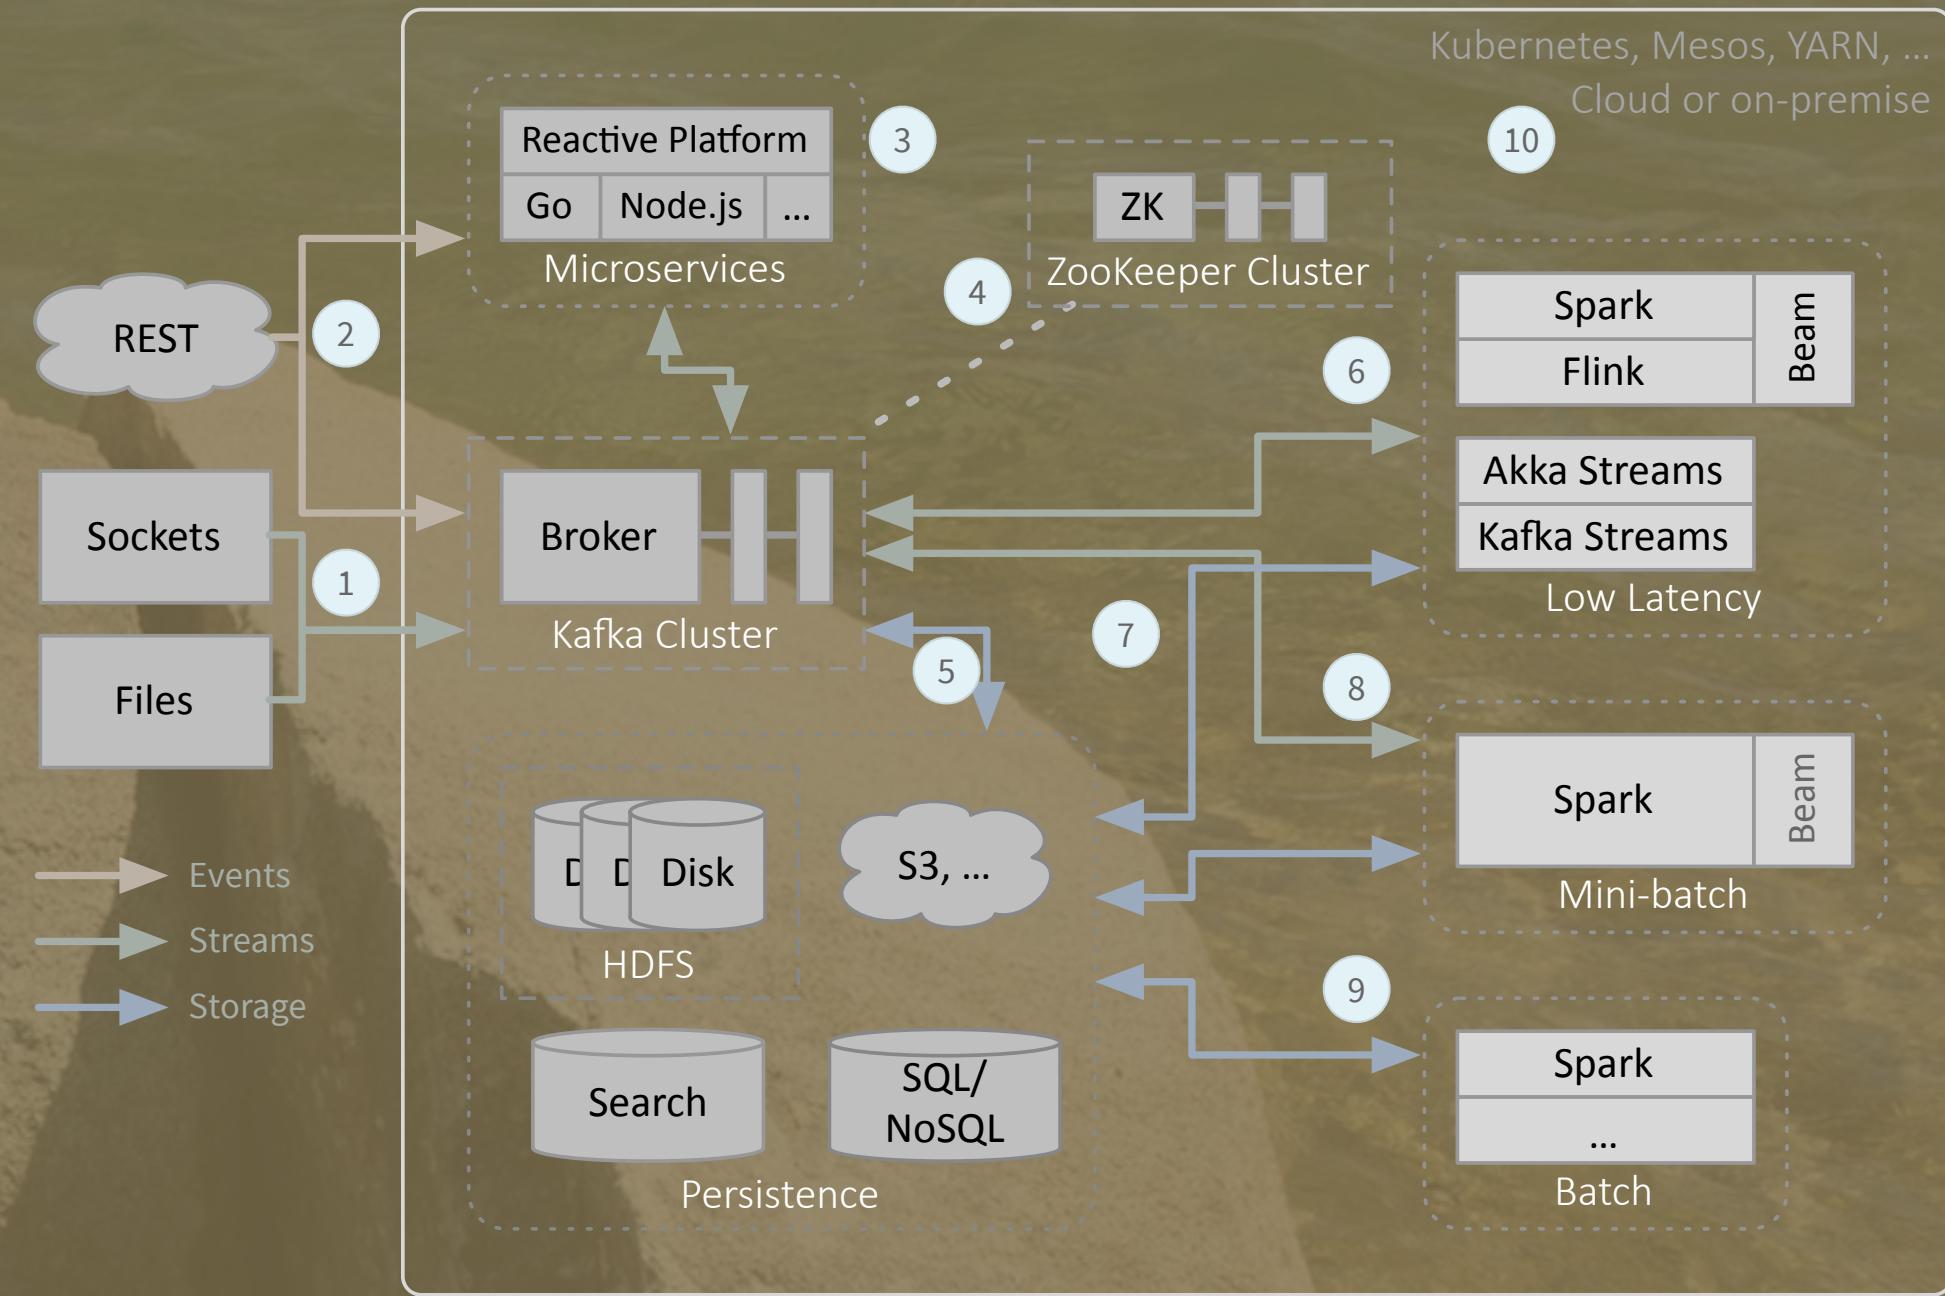
... like this?



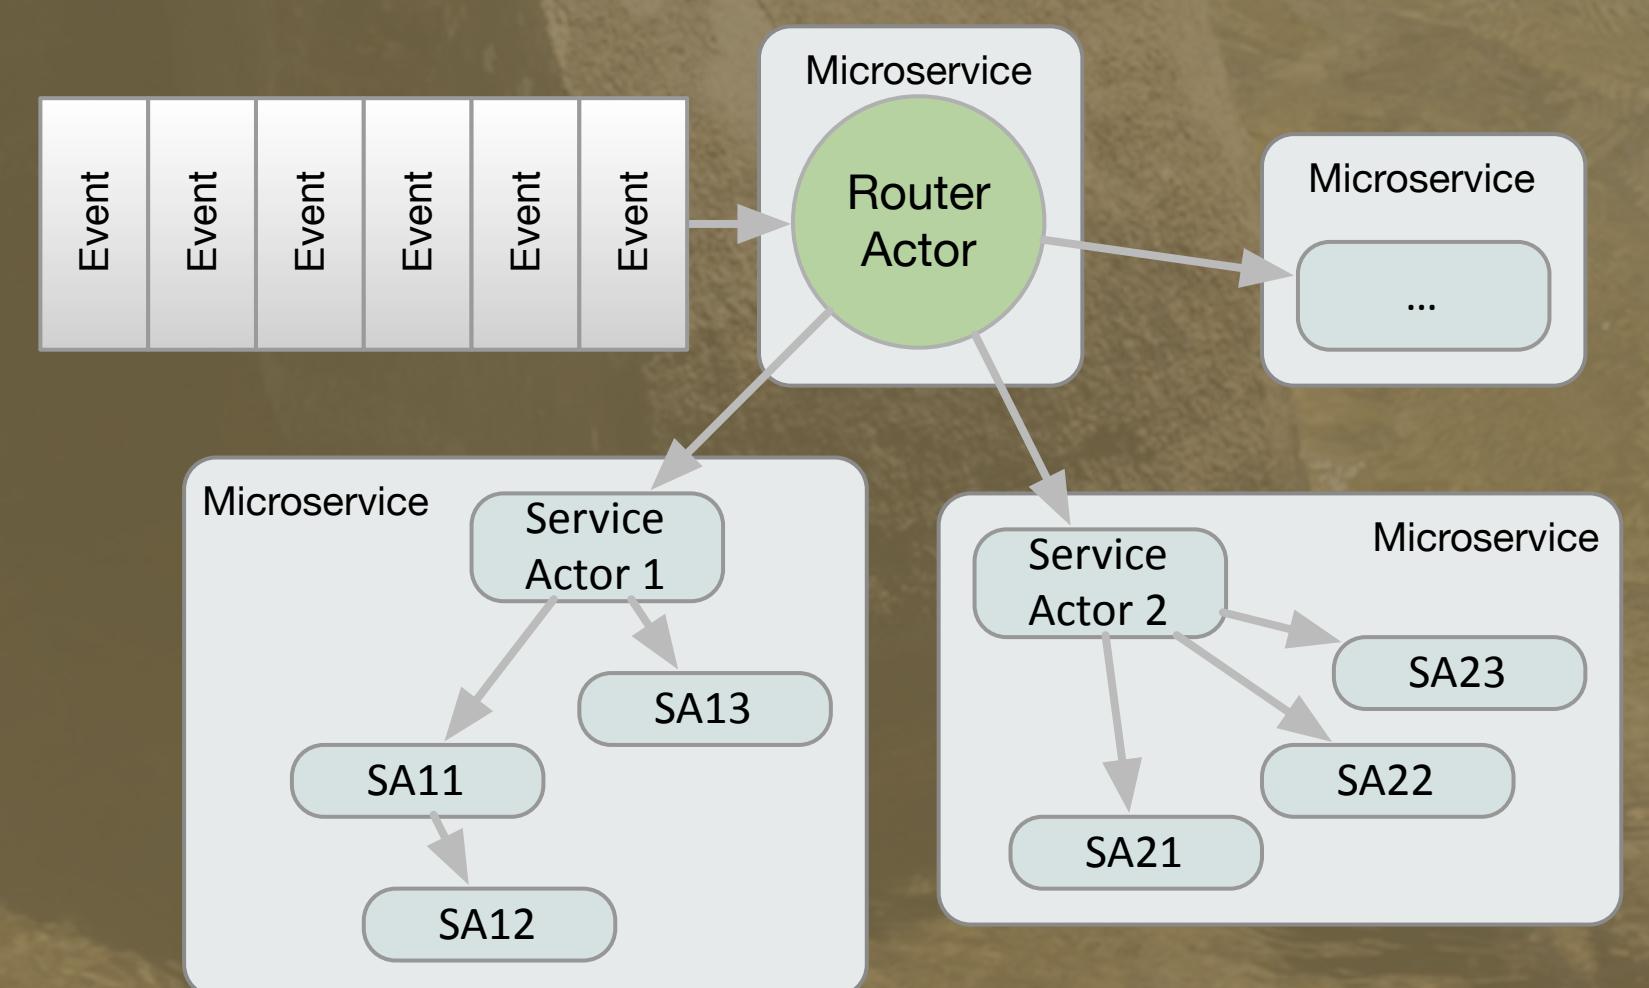
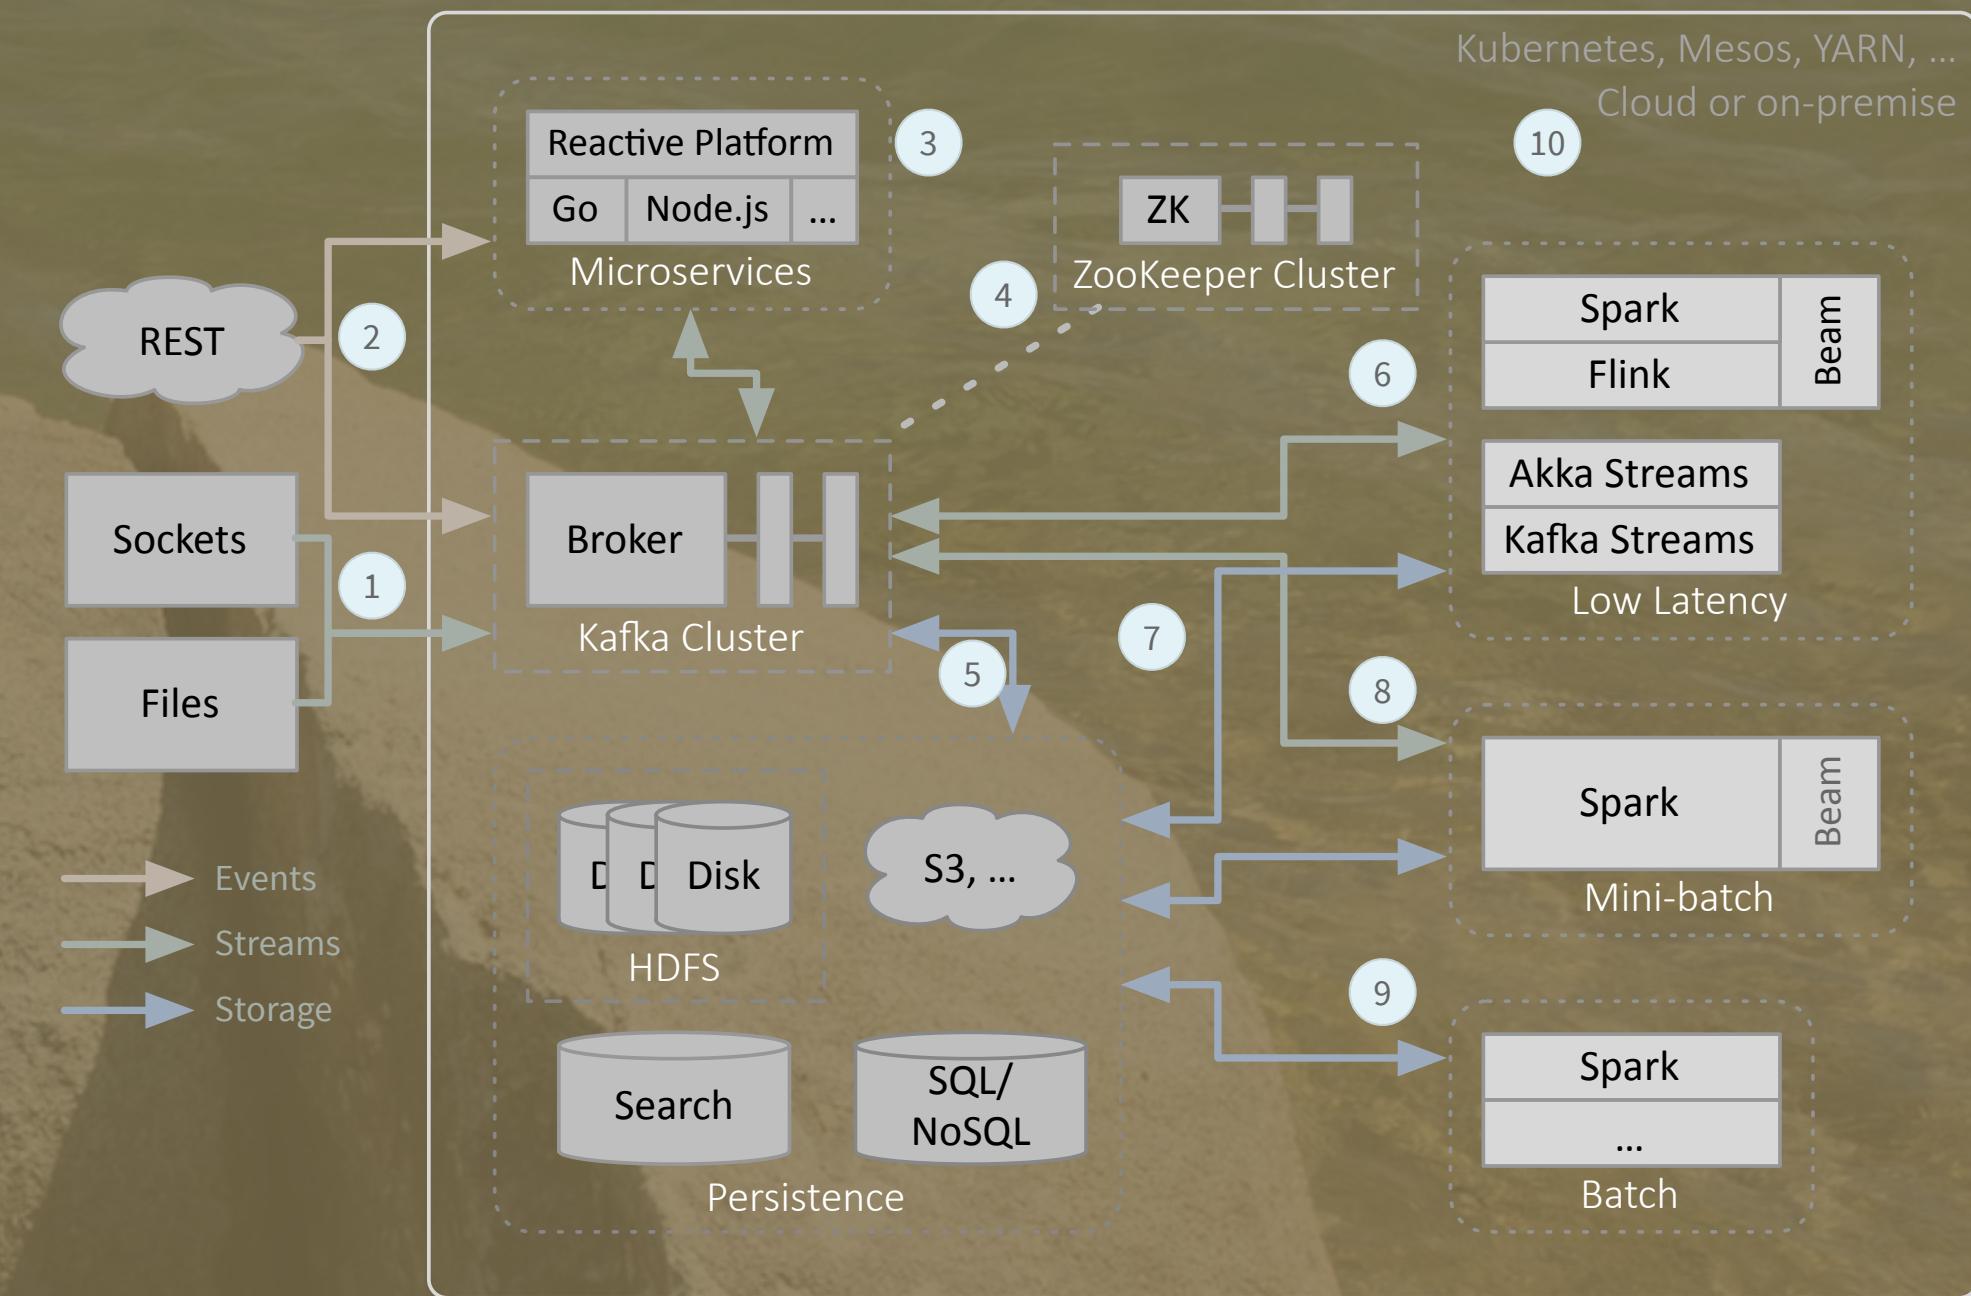
- A data app / microservice:
- A single responsibility.
- ...



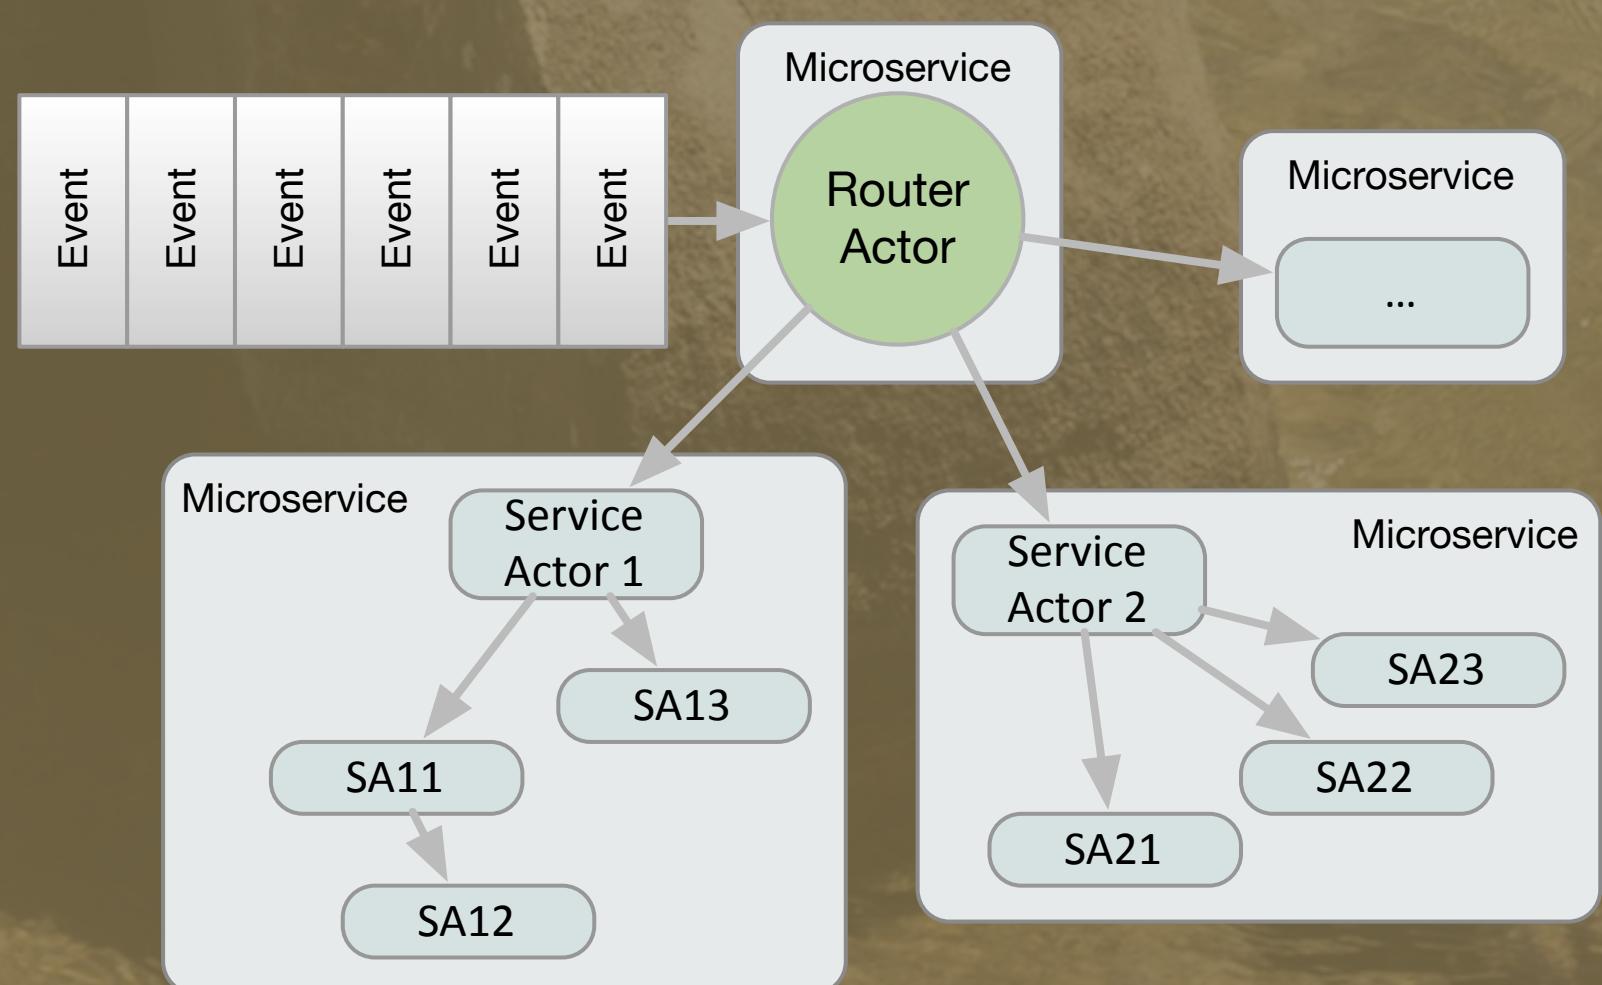
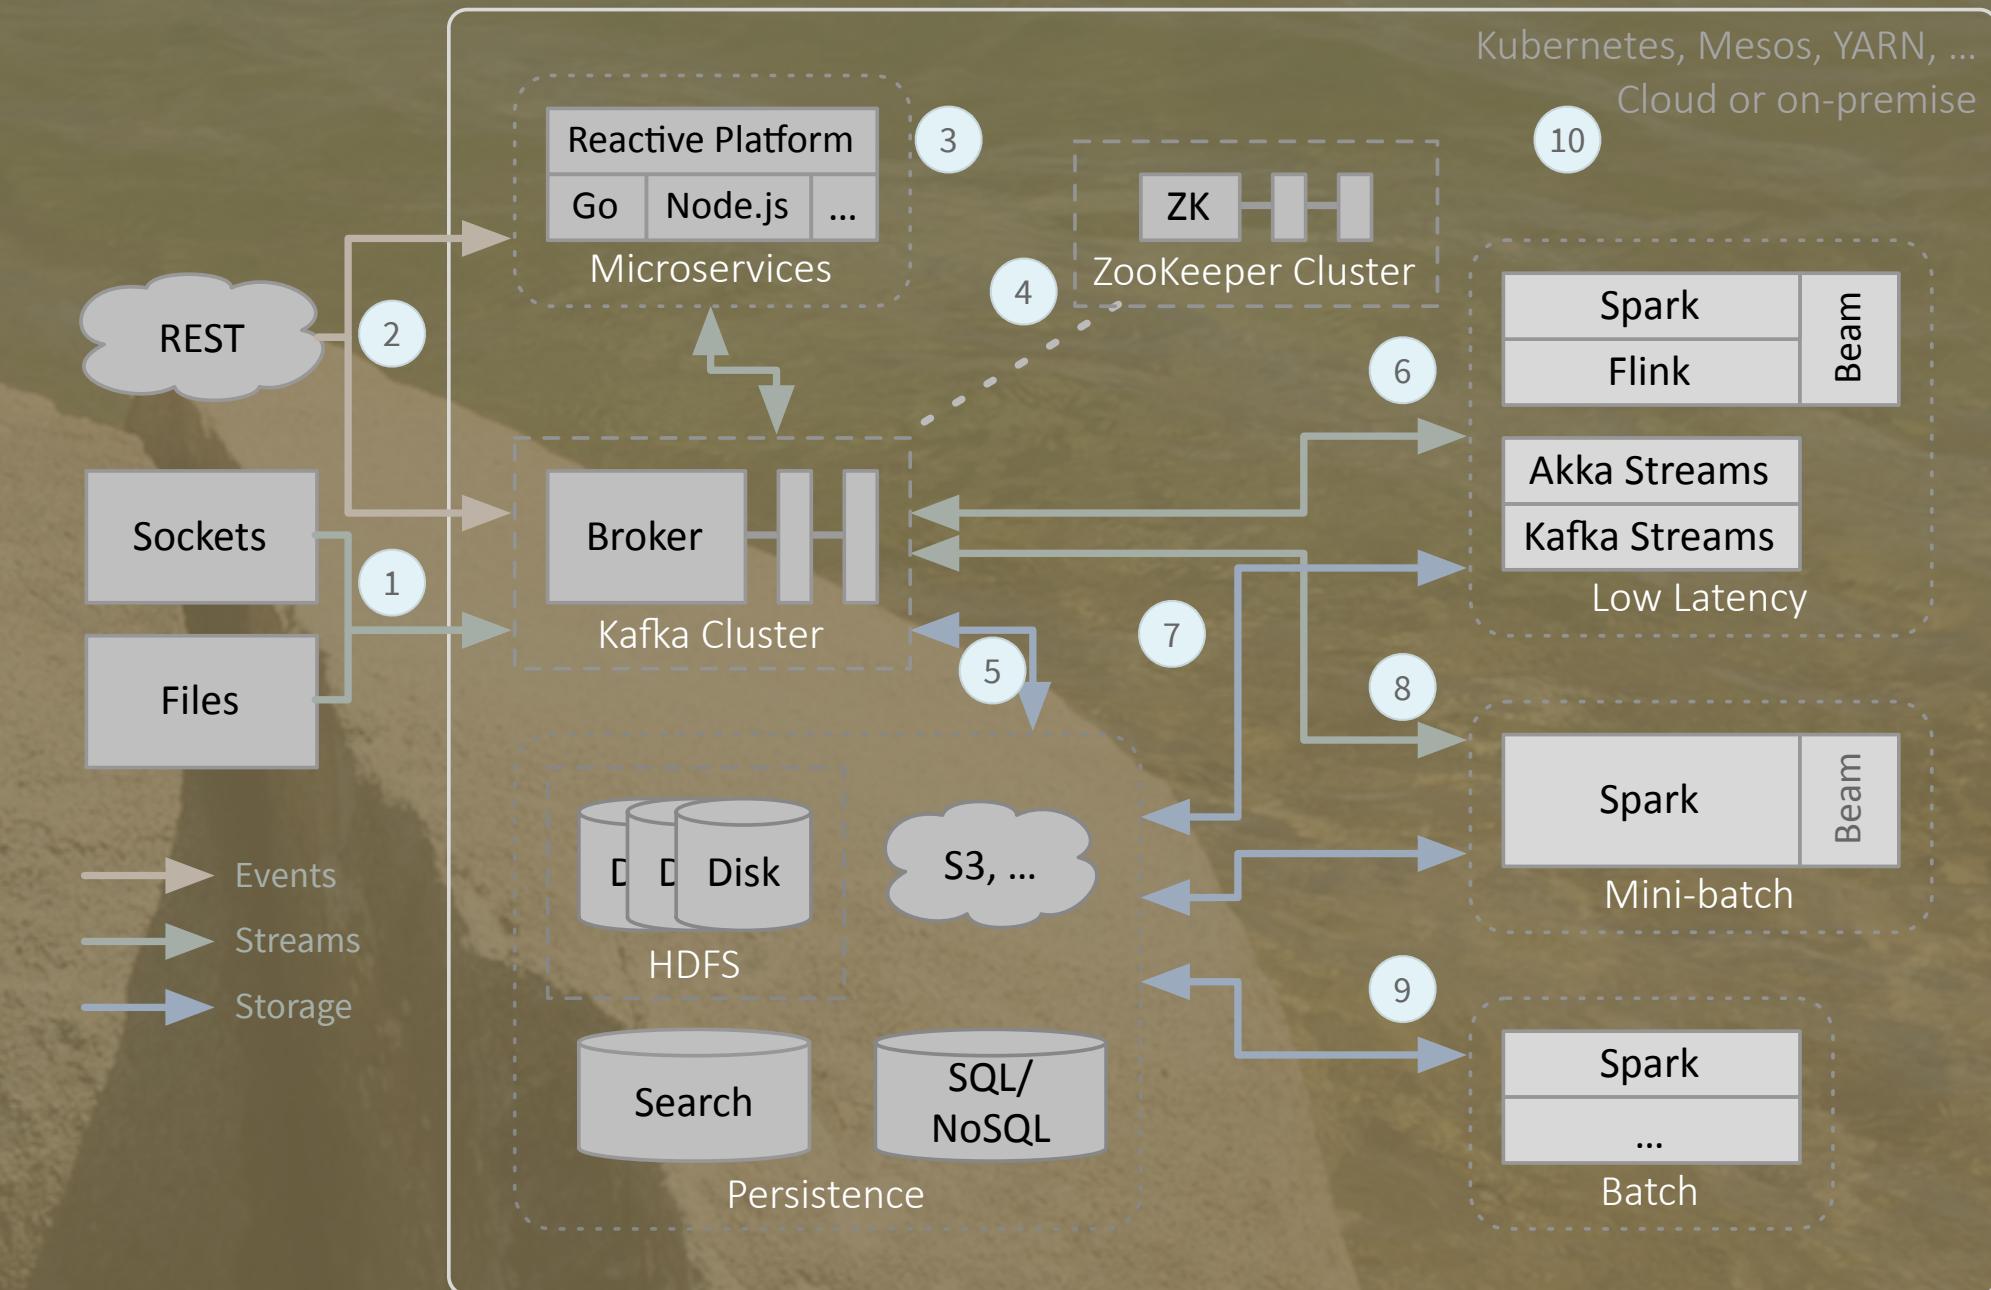
- A data app / microservice:
- A single responsibility.
- The input never ends.



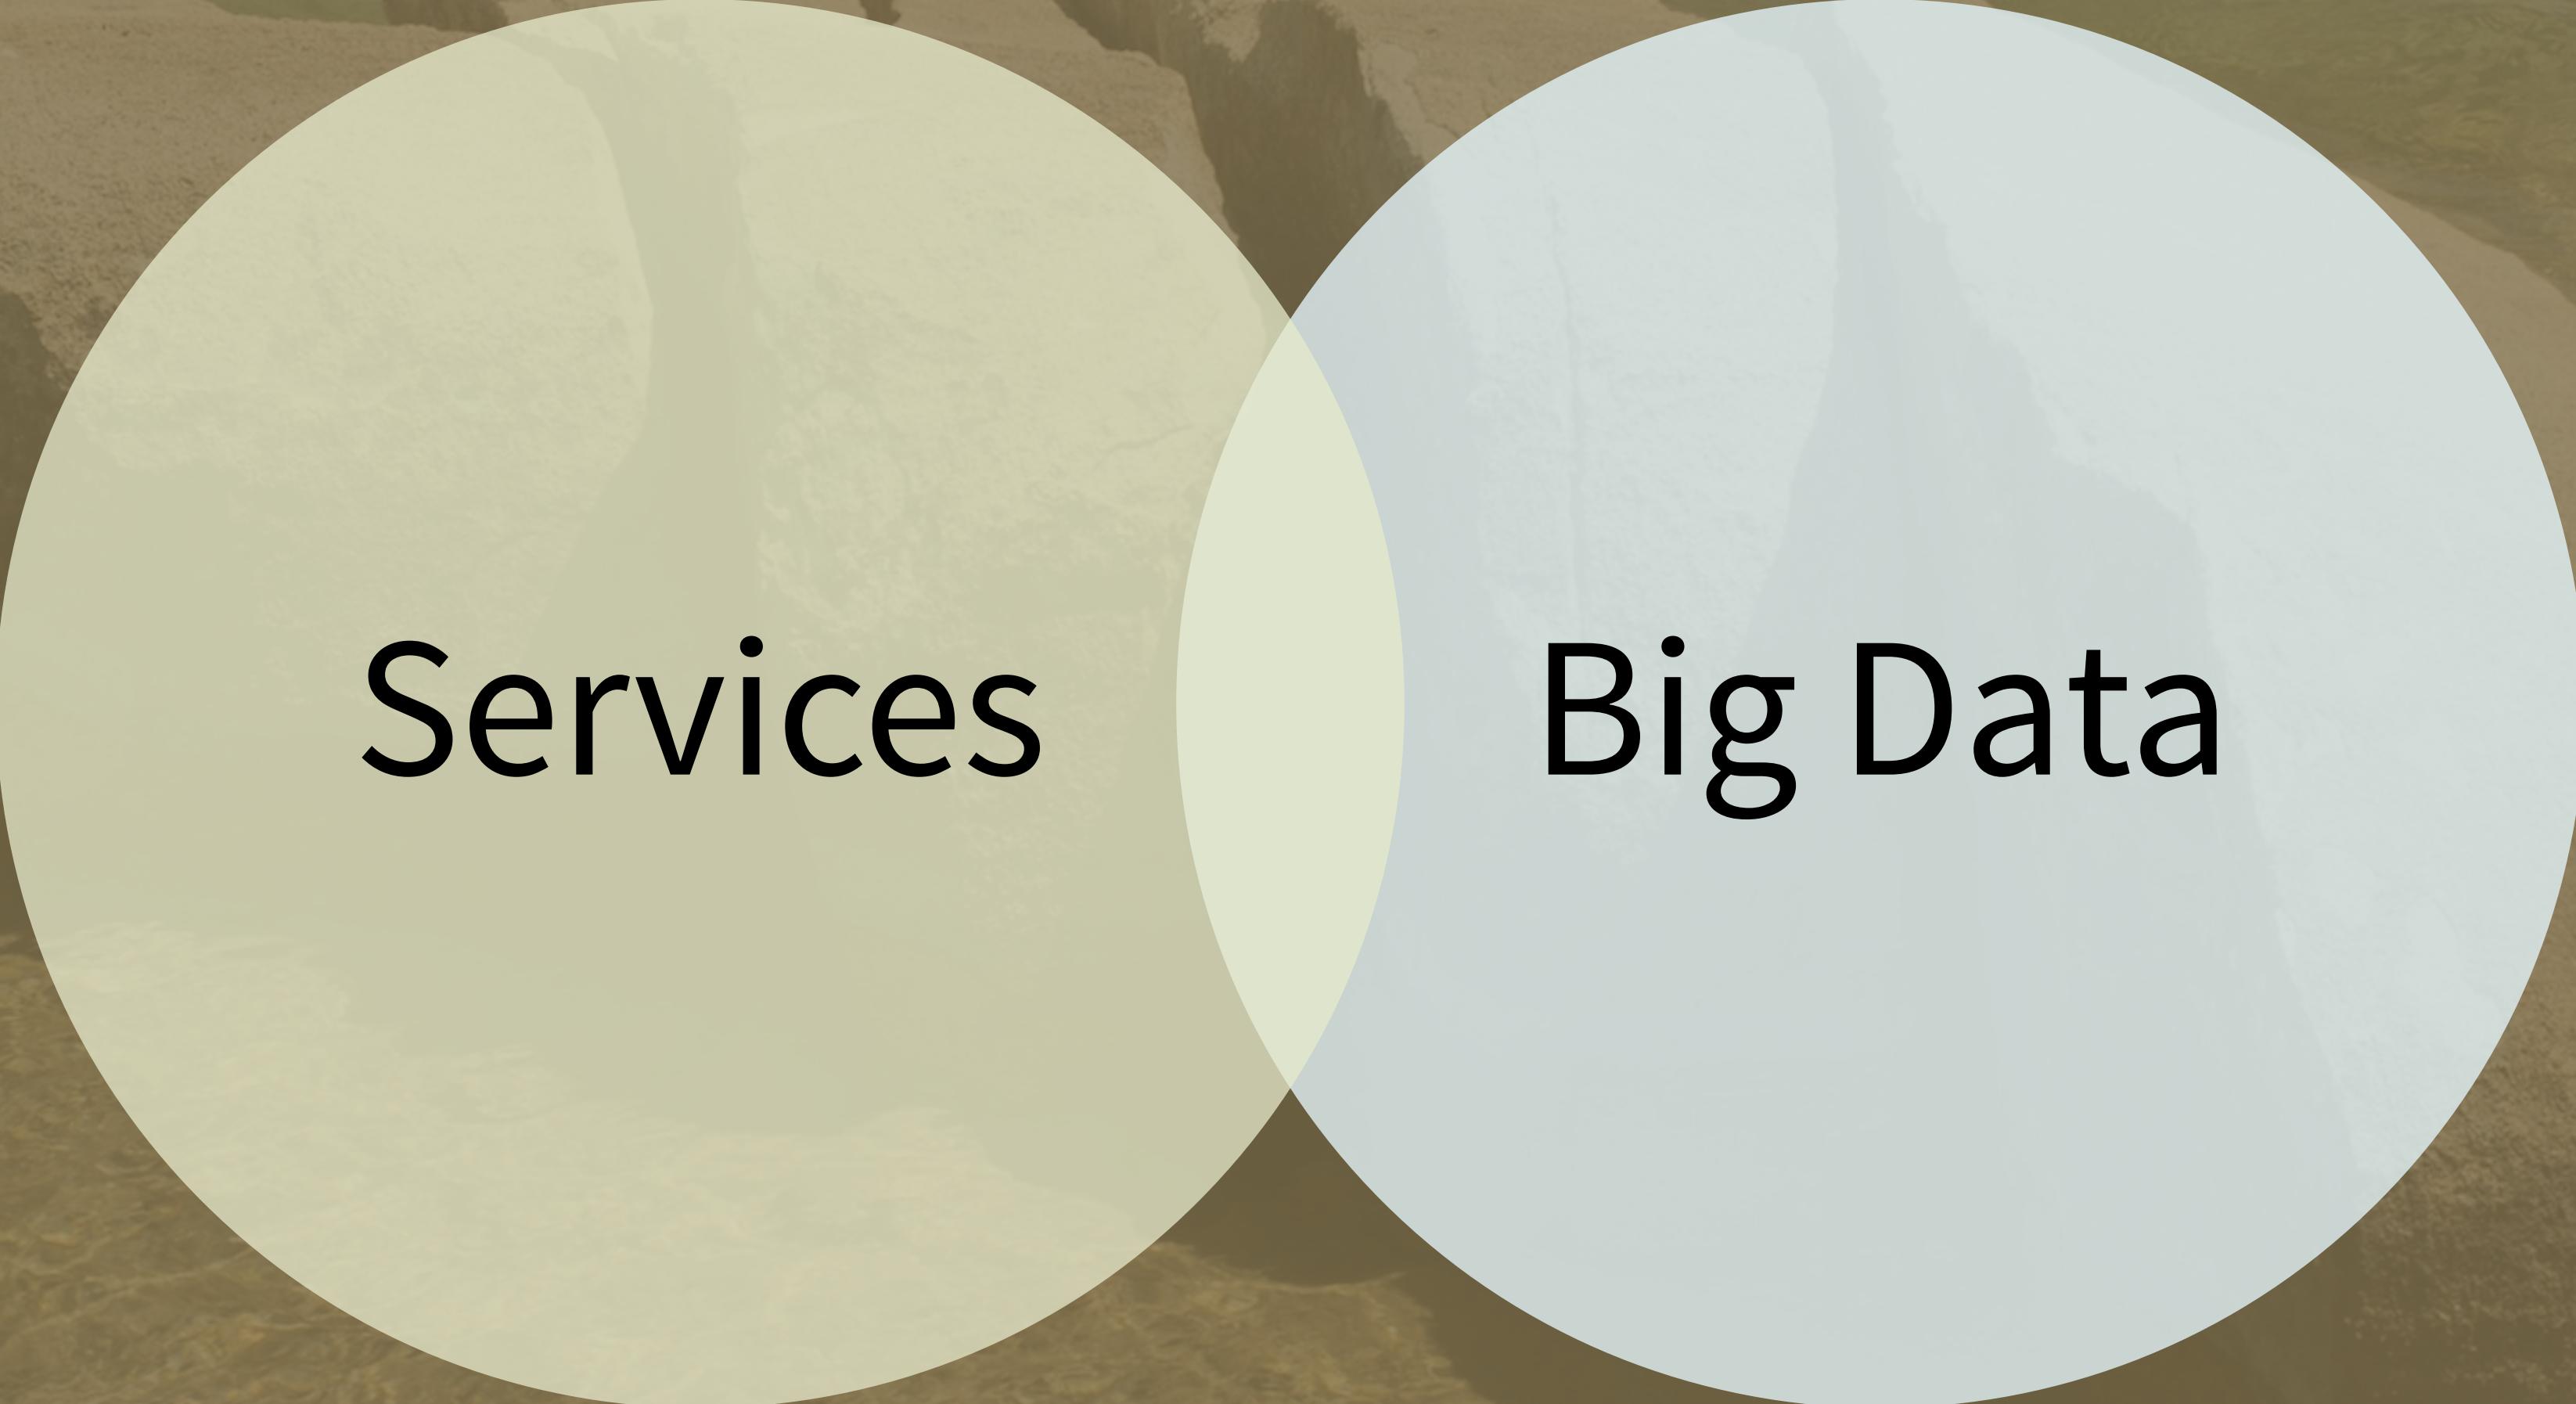
- A data app/microservice:
- A single responsibility.
- The input never ends.
- So, both must be available, responsive, resilient, & scalable. I.e., reactive



- Going the other way, “small” microservice architectures become data-centric, as the data grows.



The Recent Past



Services

Big Data

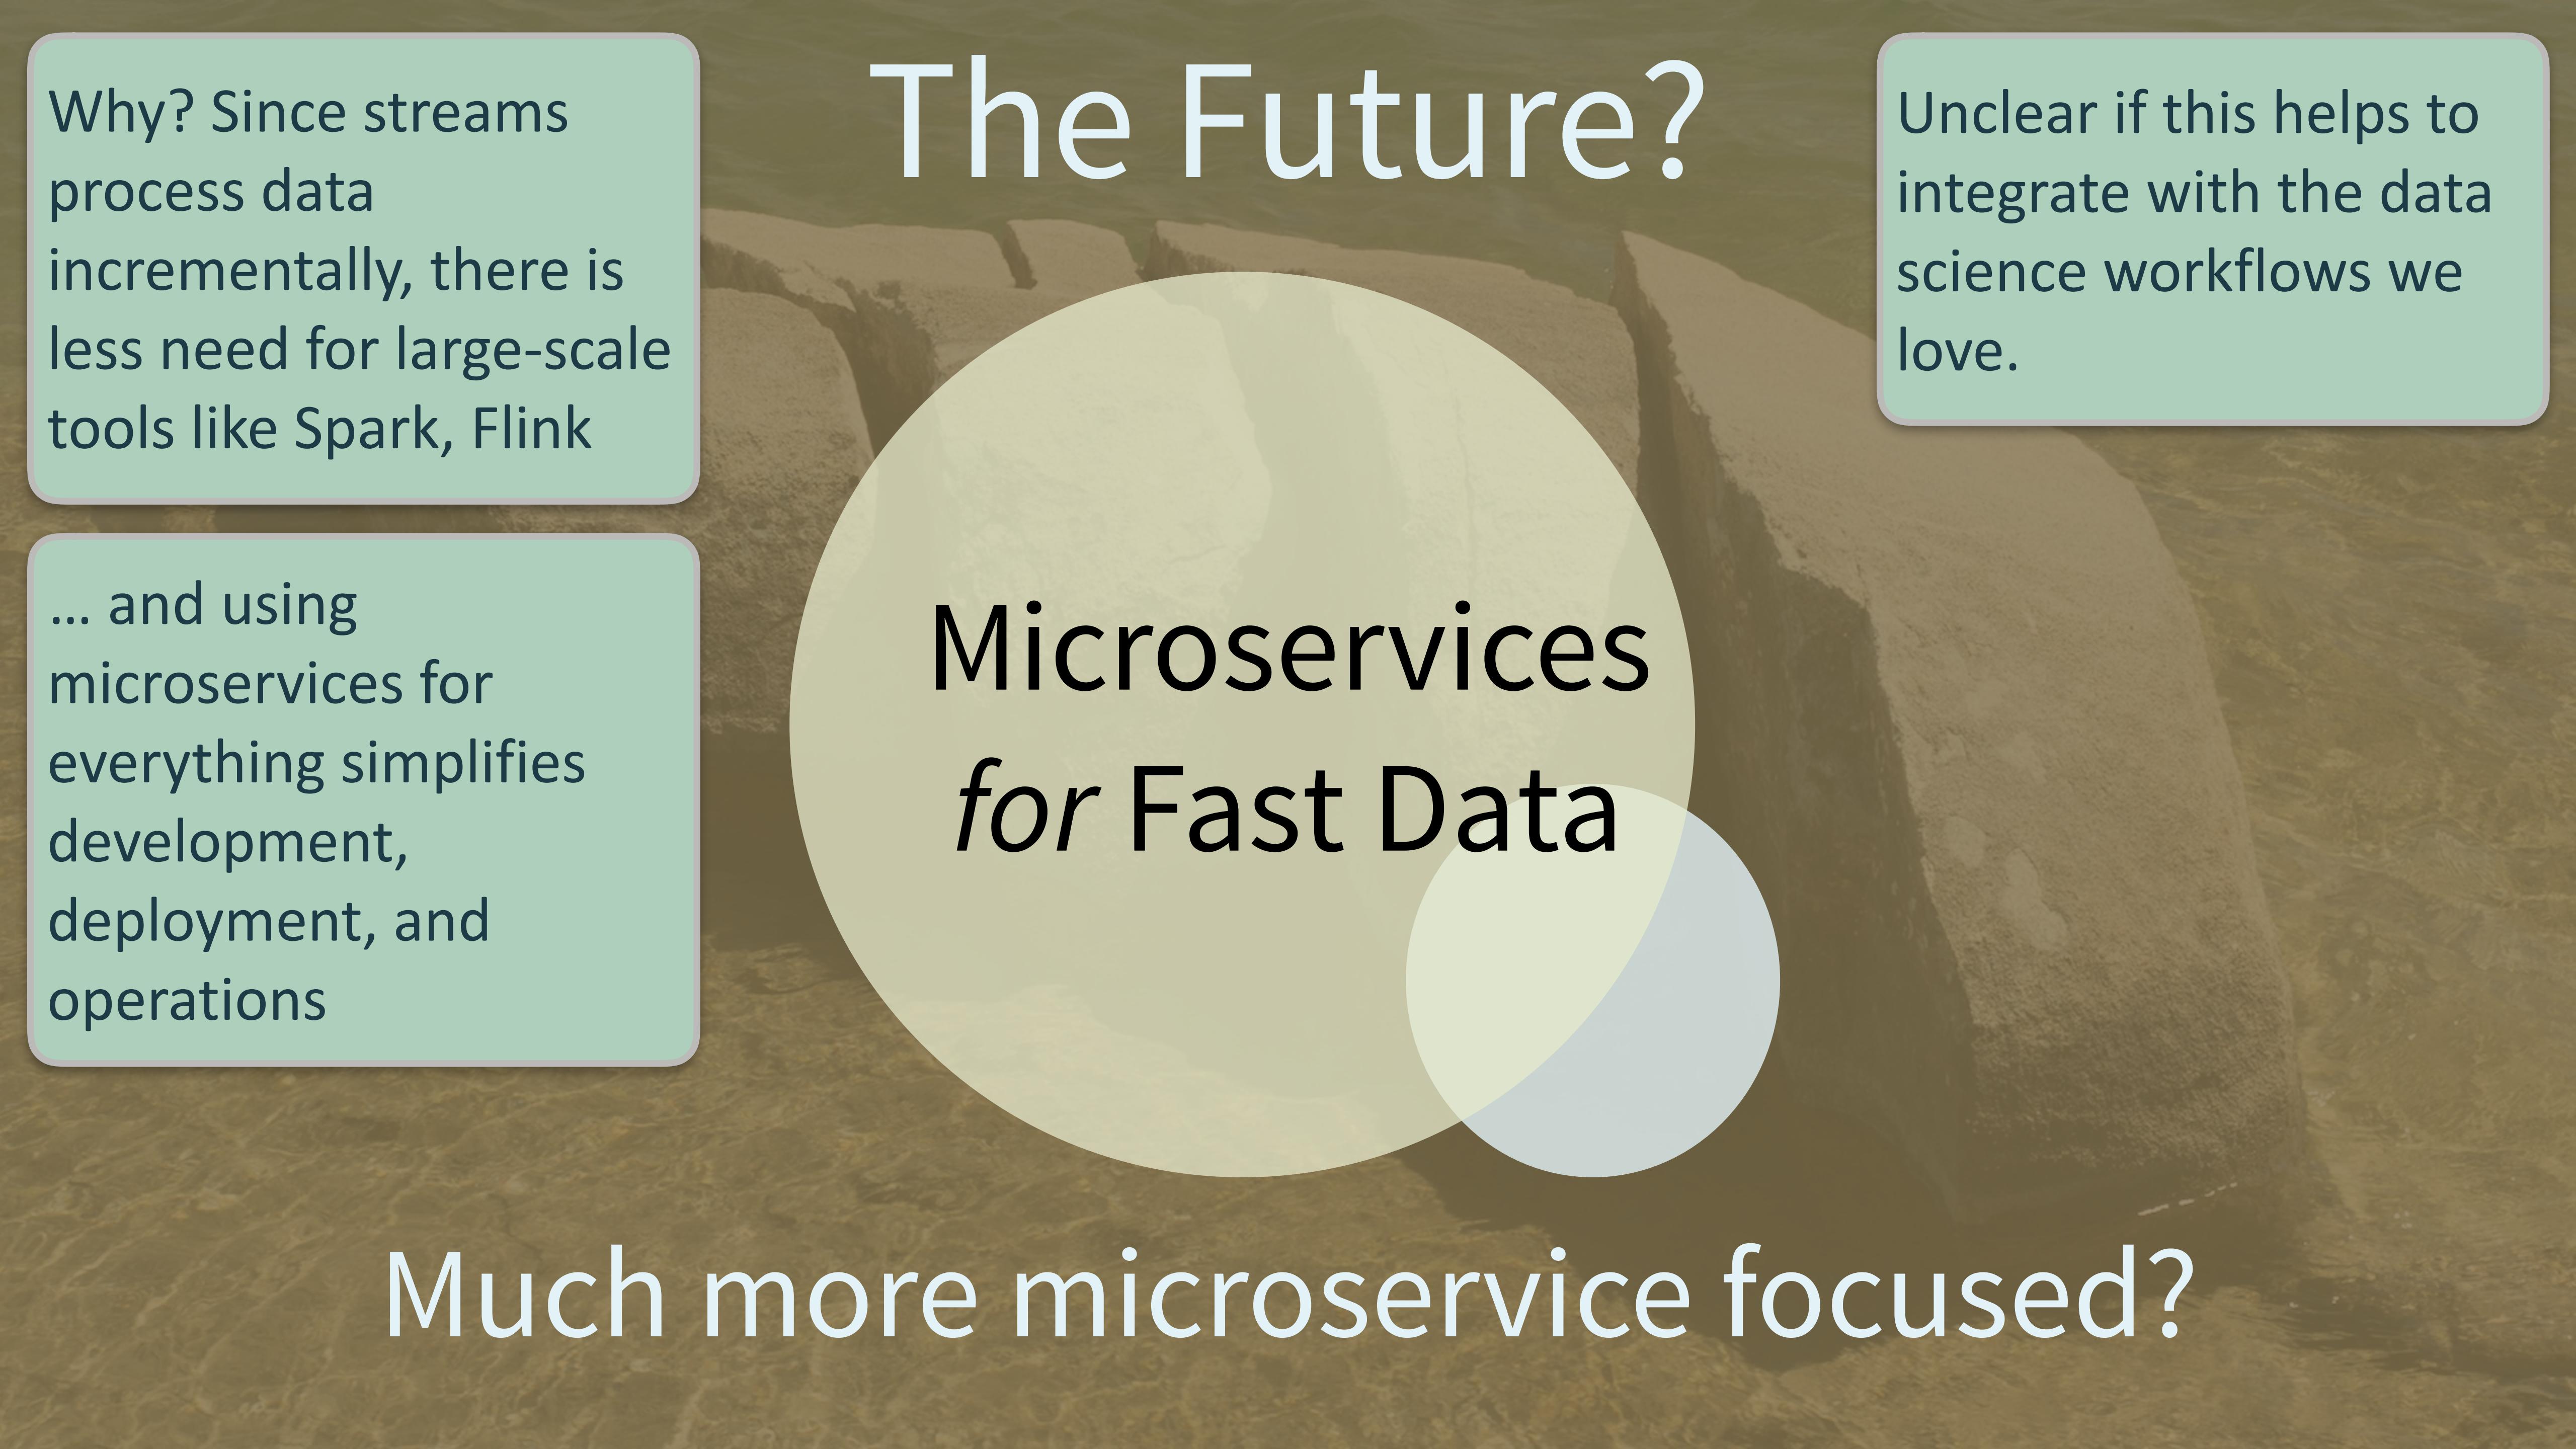
Some Overlap: Concerns, Architecture

The Present



Microservices
& *Fast* Data

Much More Overlap



Why? Since streams process data incrementally, there is less need for large-scale tools like Spark, Flink

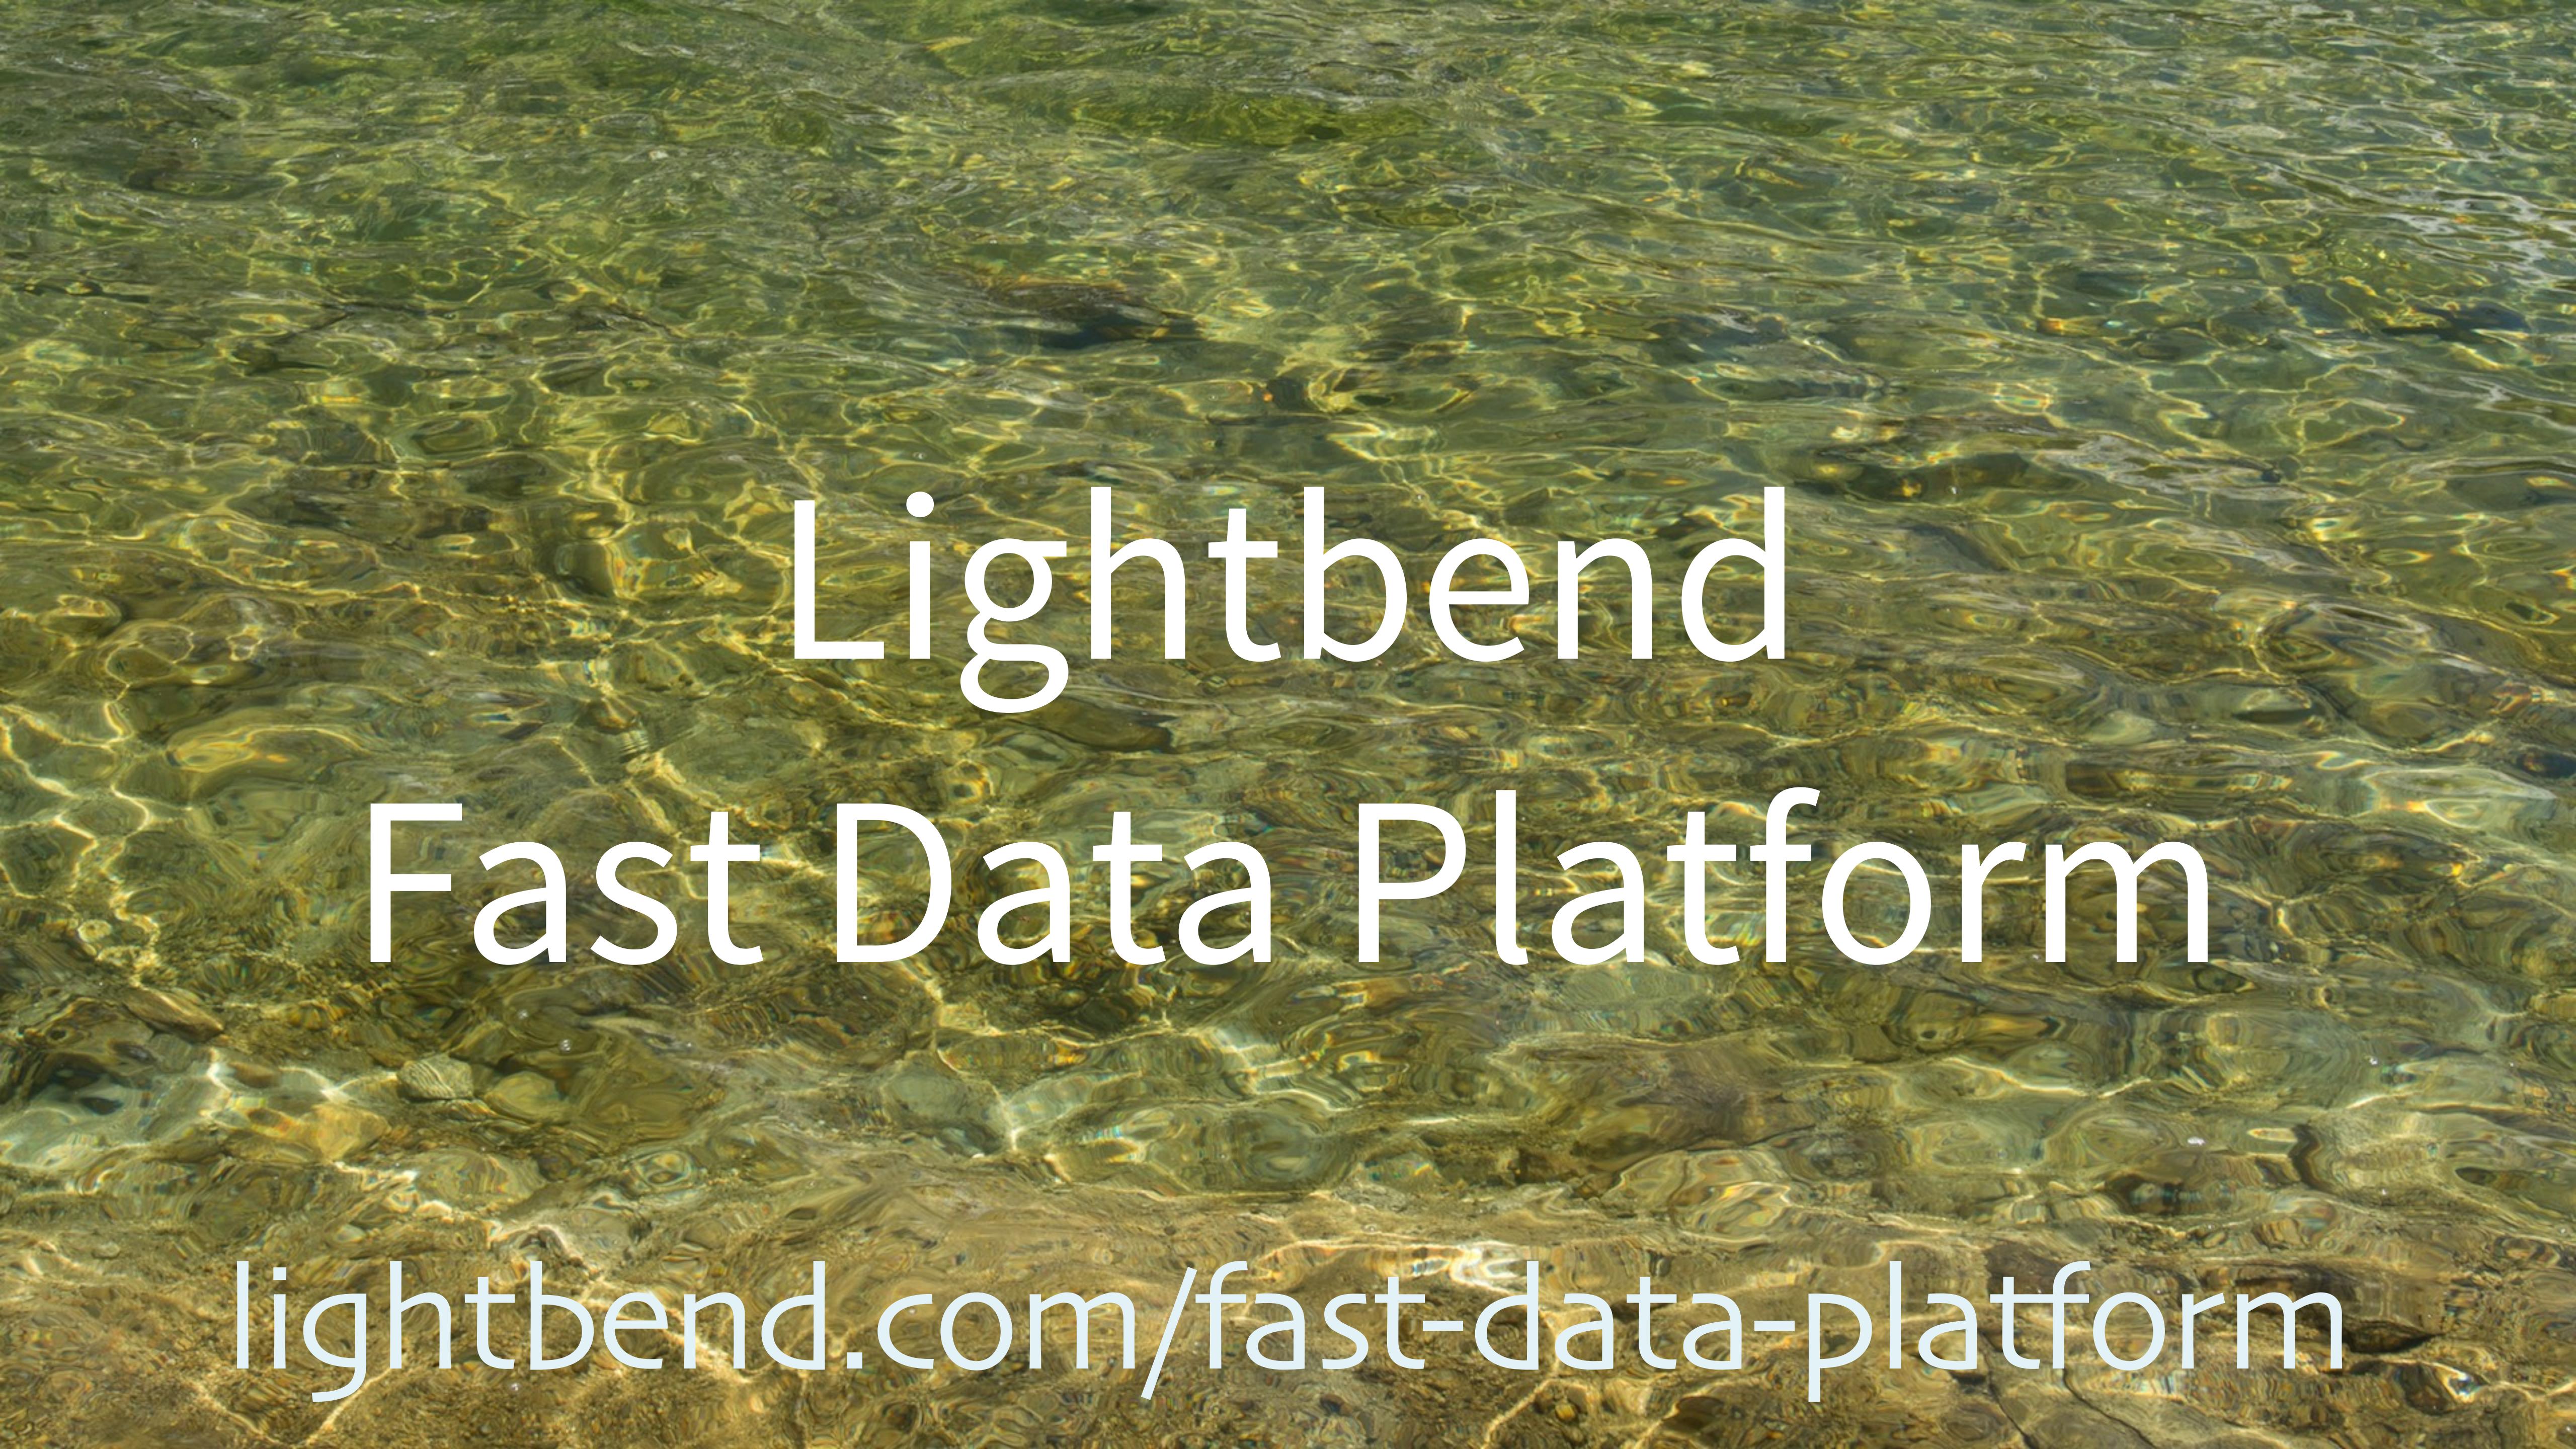
... and using microservices for everything simplifies development, deployment, and operations

The Future?

Unclear if this helps to integrate with the data science workflows we love.

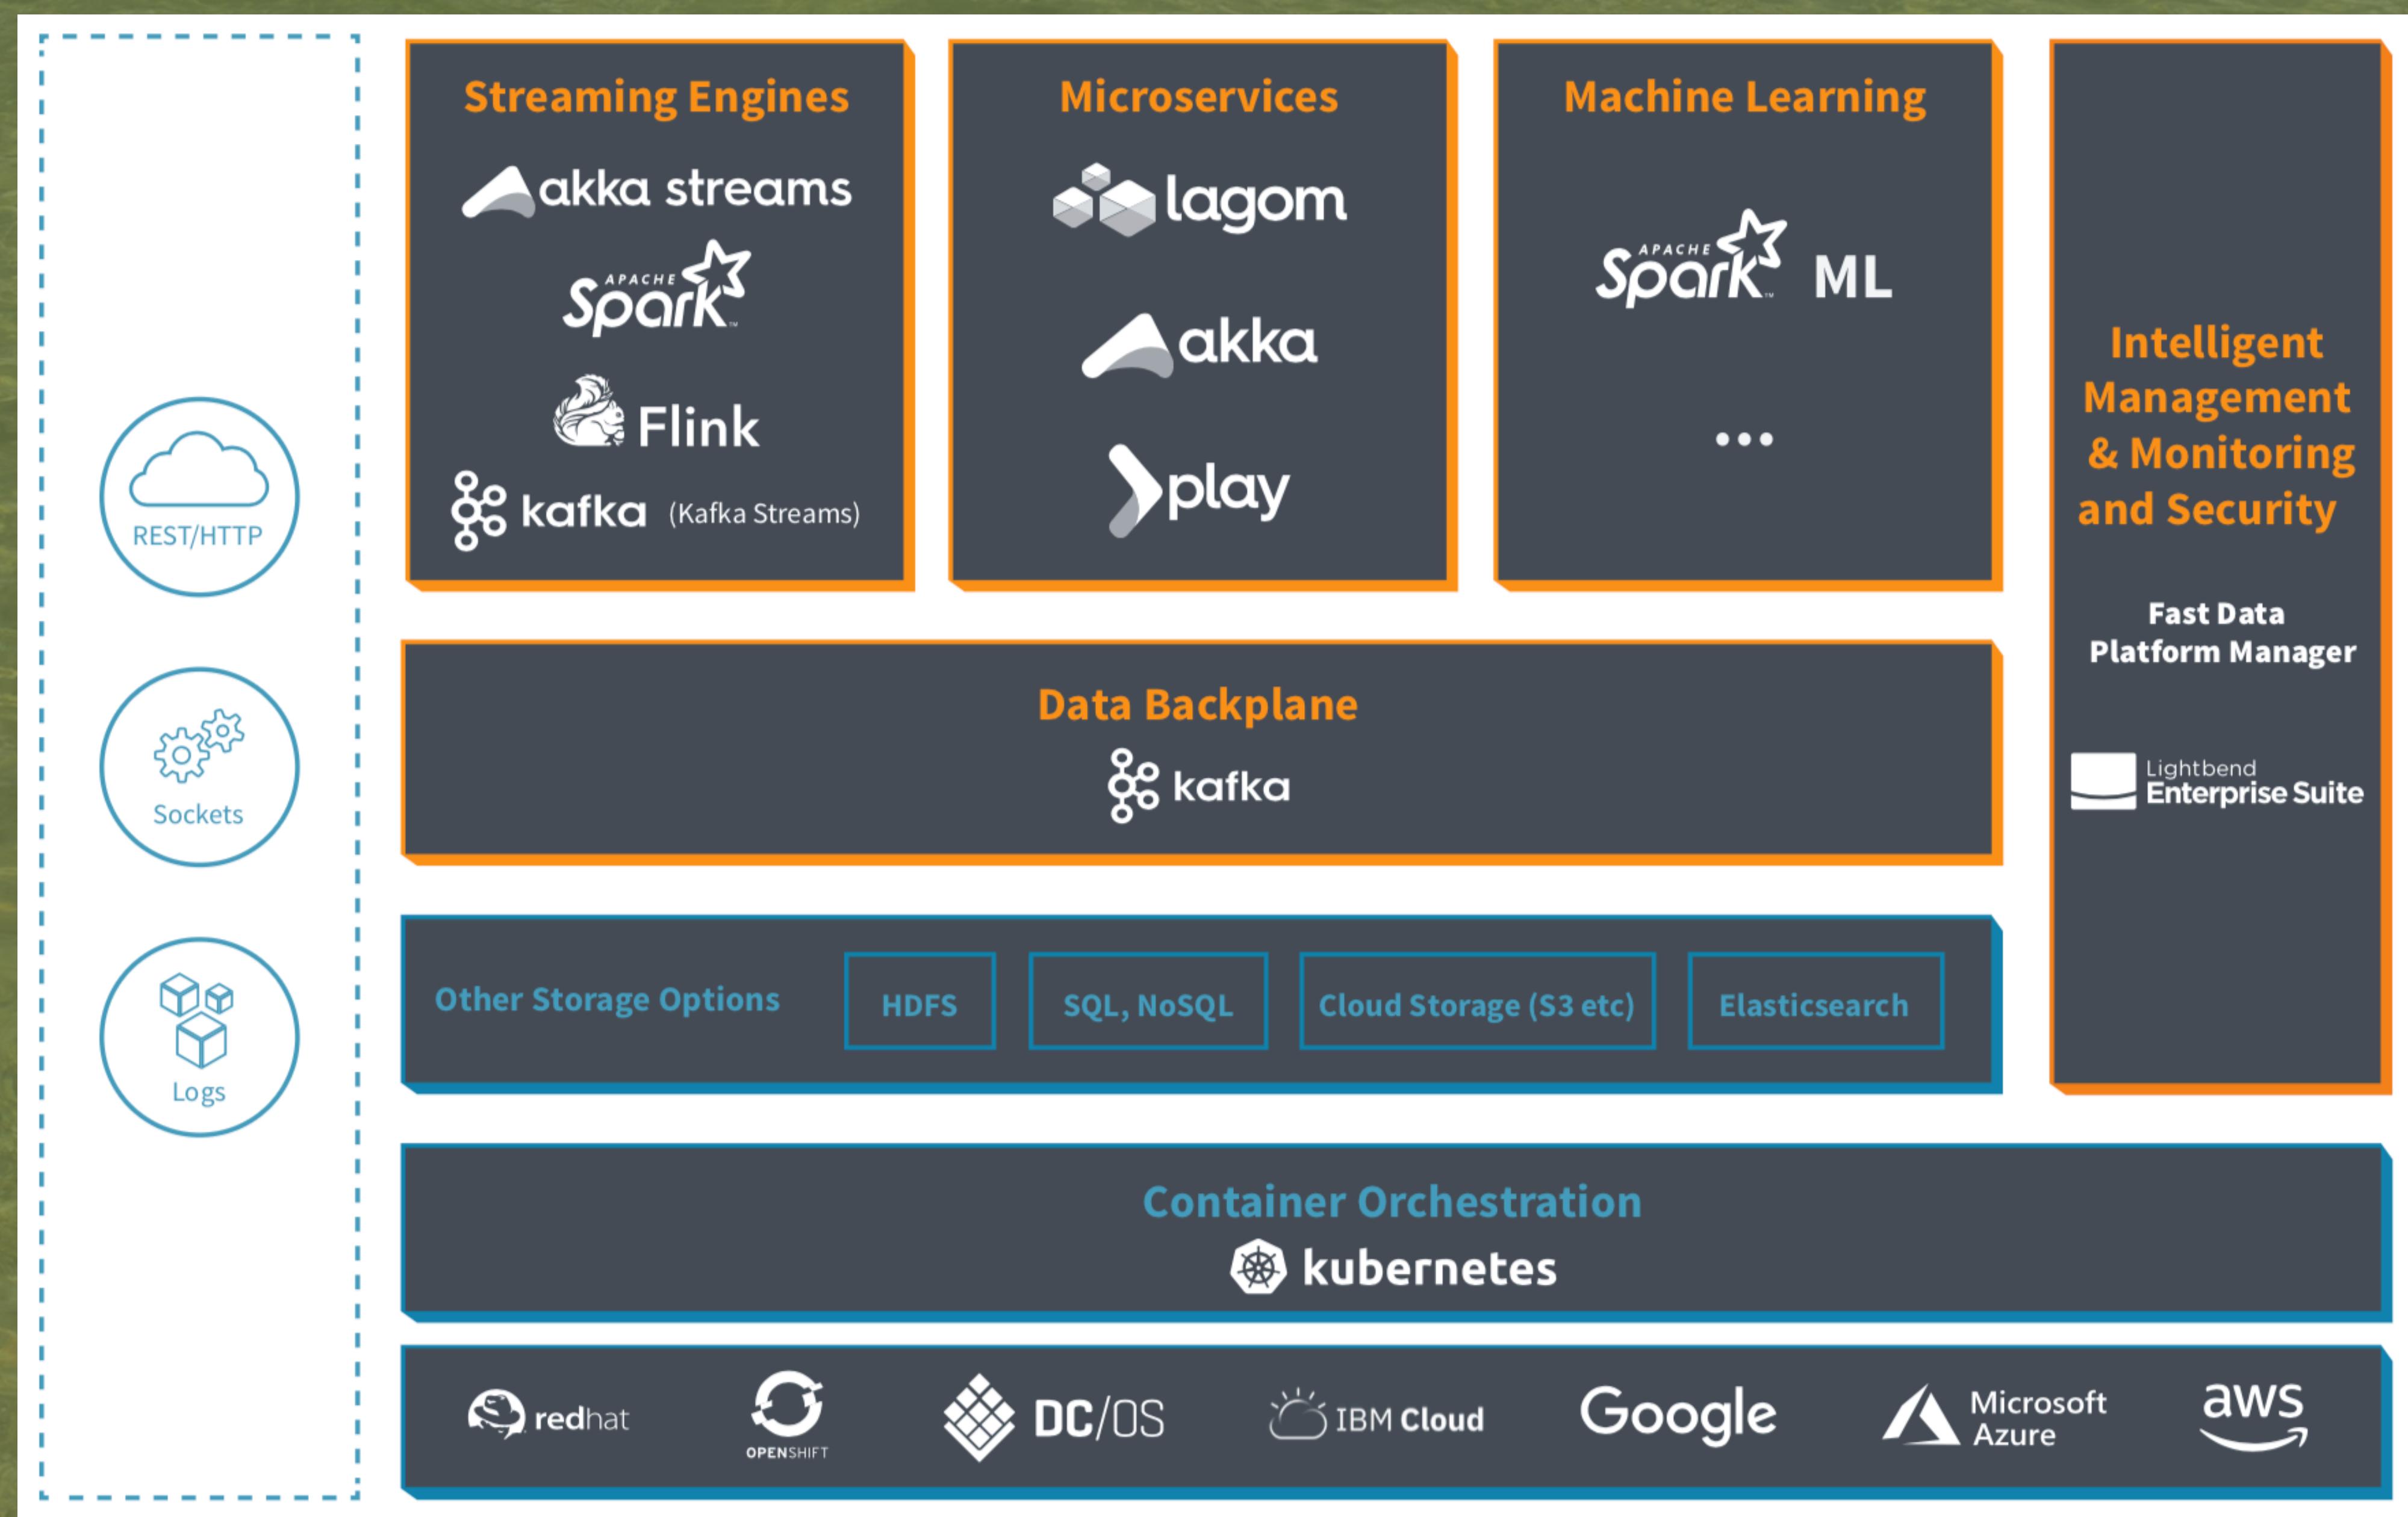
Microservices
for Fast Data

Much more microservice focused?



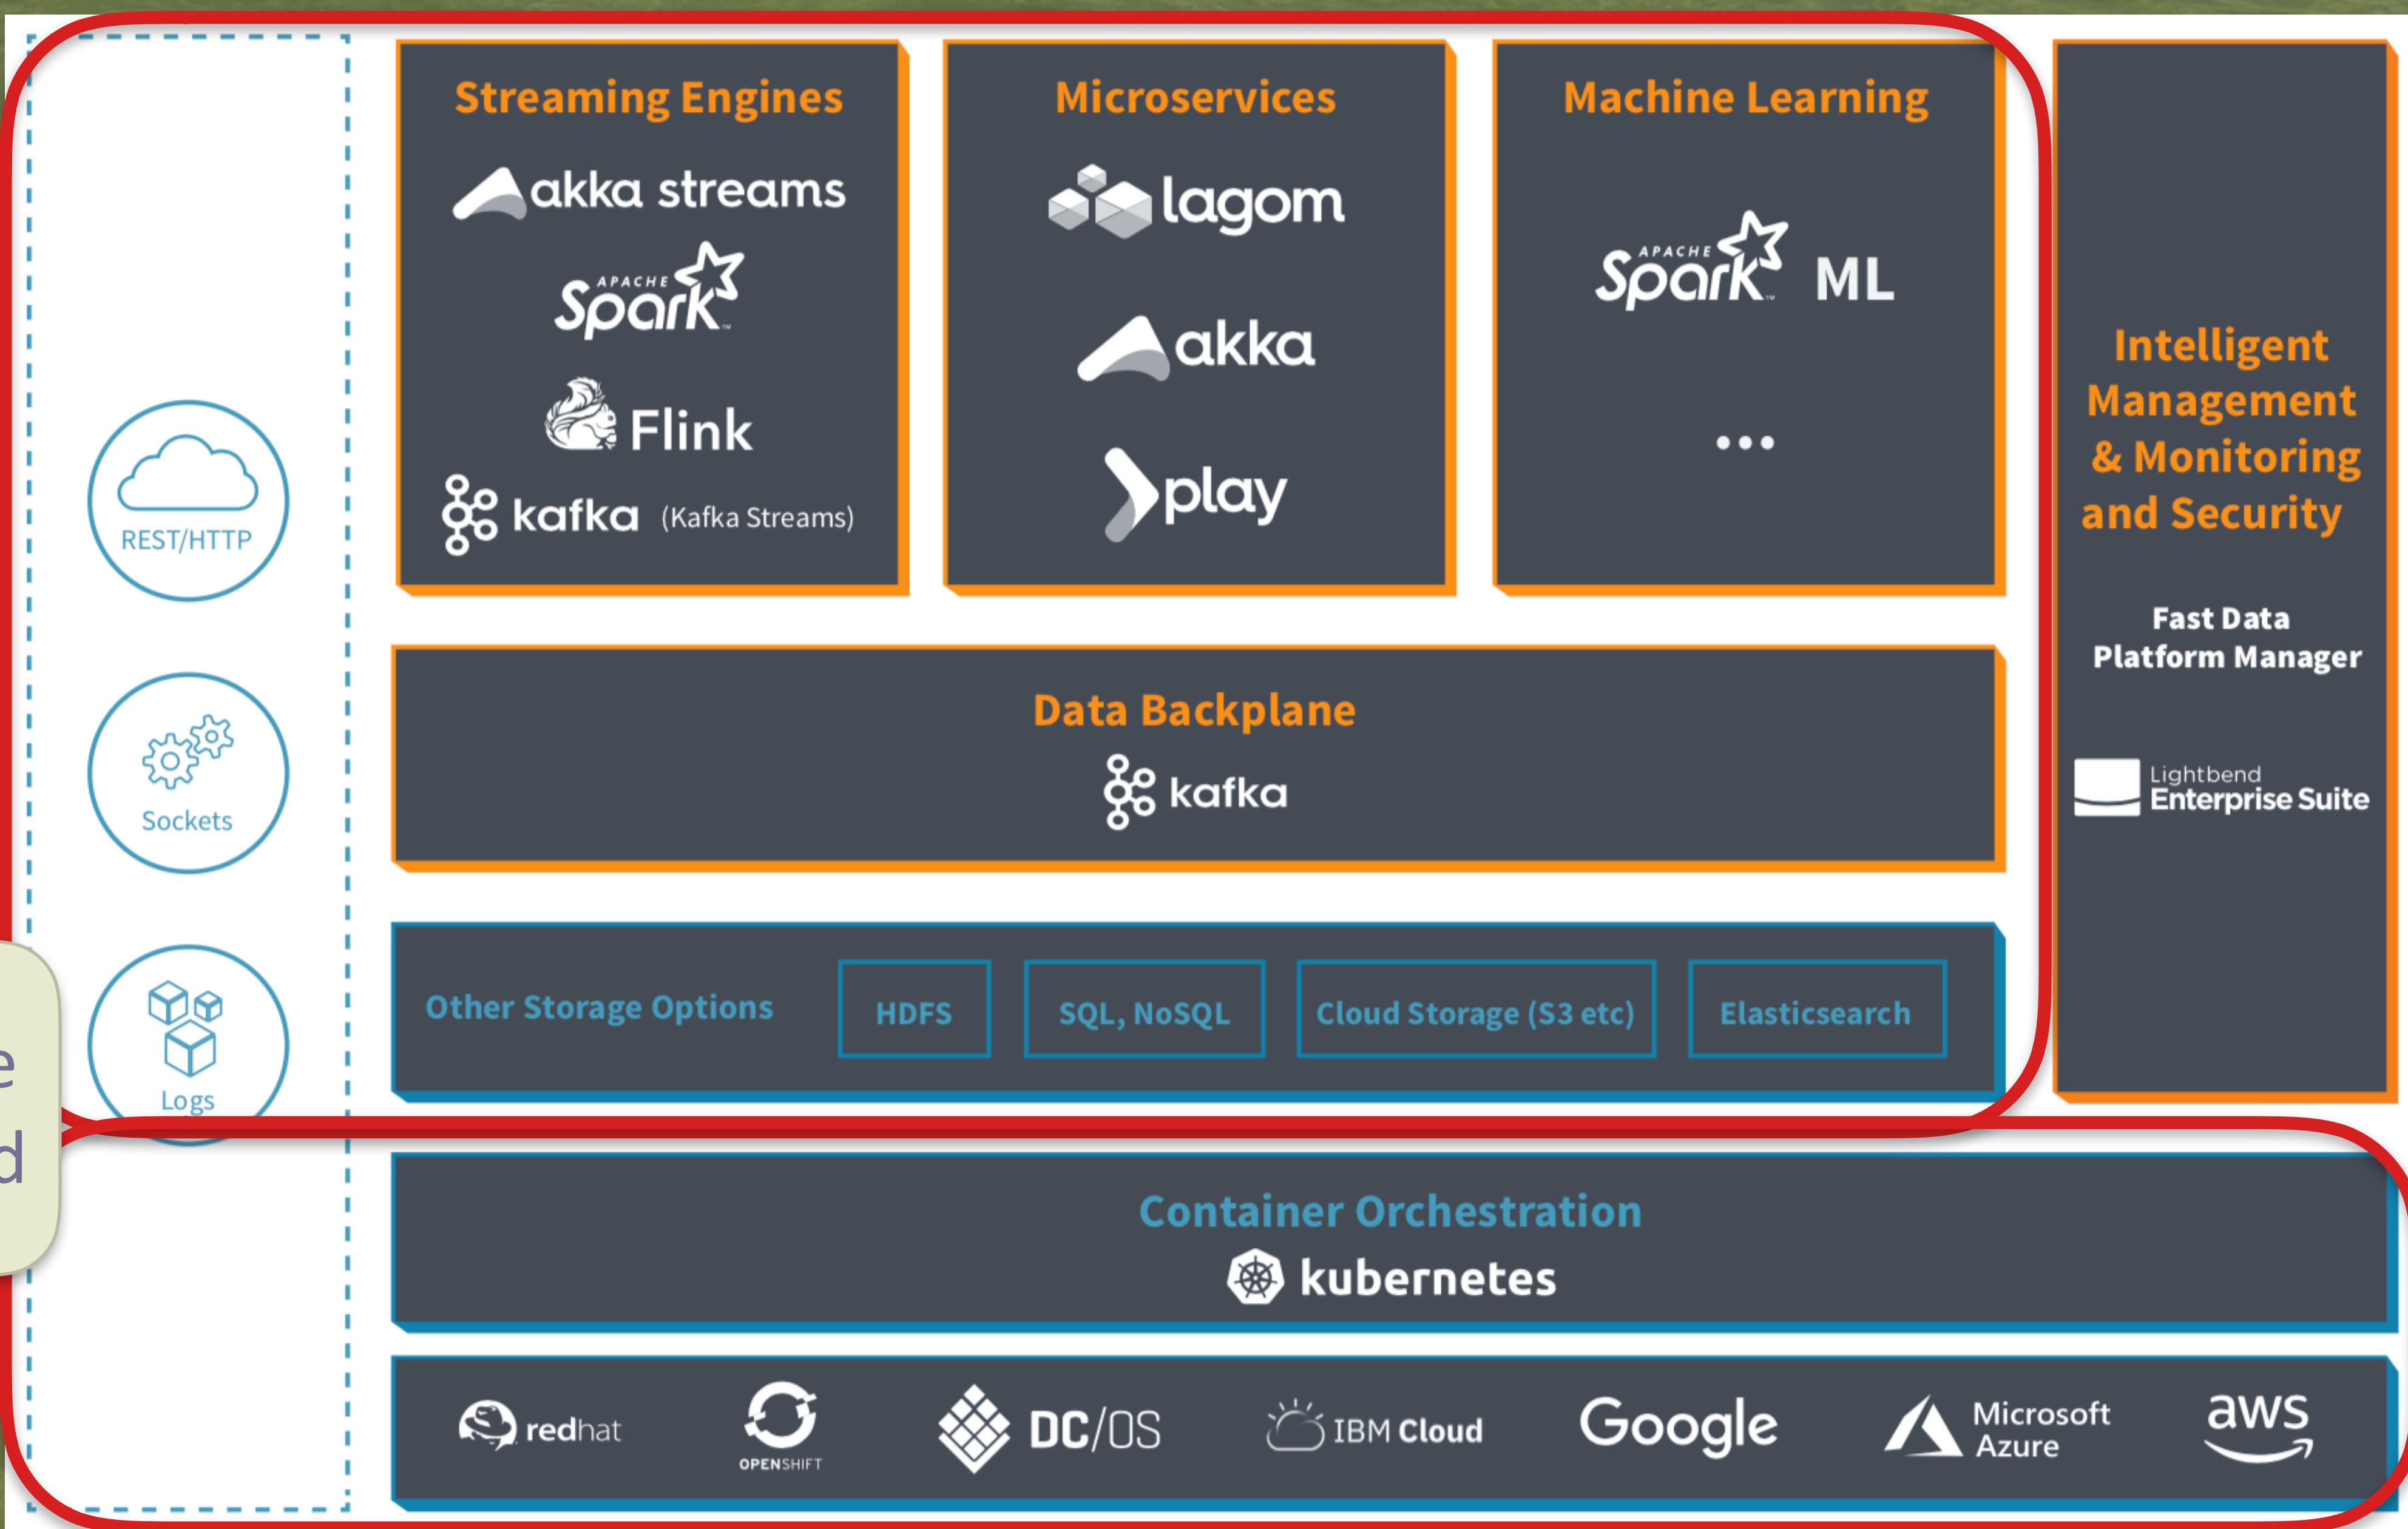
Lightbend Fast Data Platform

lightbend.com/fast-data-platform

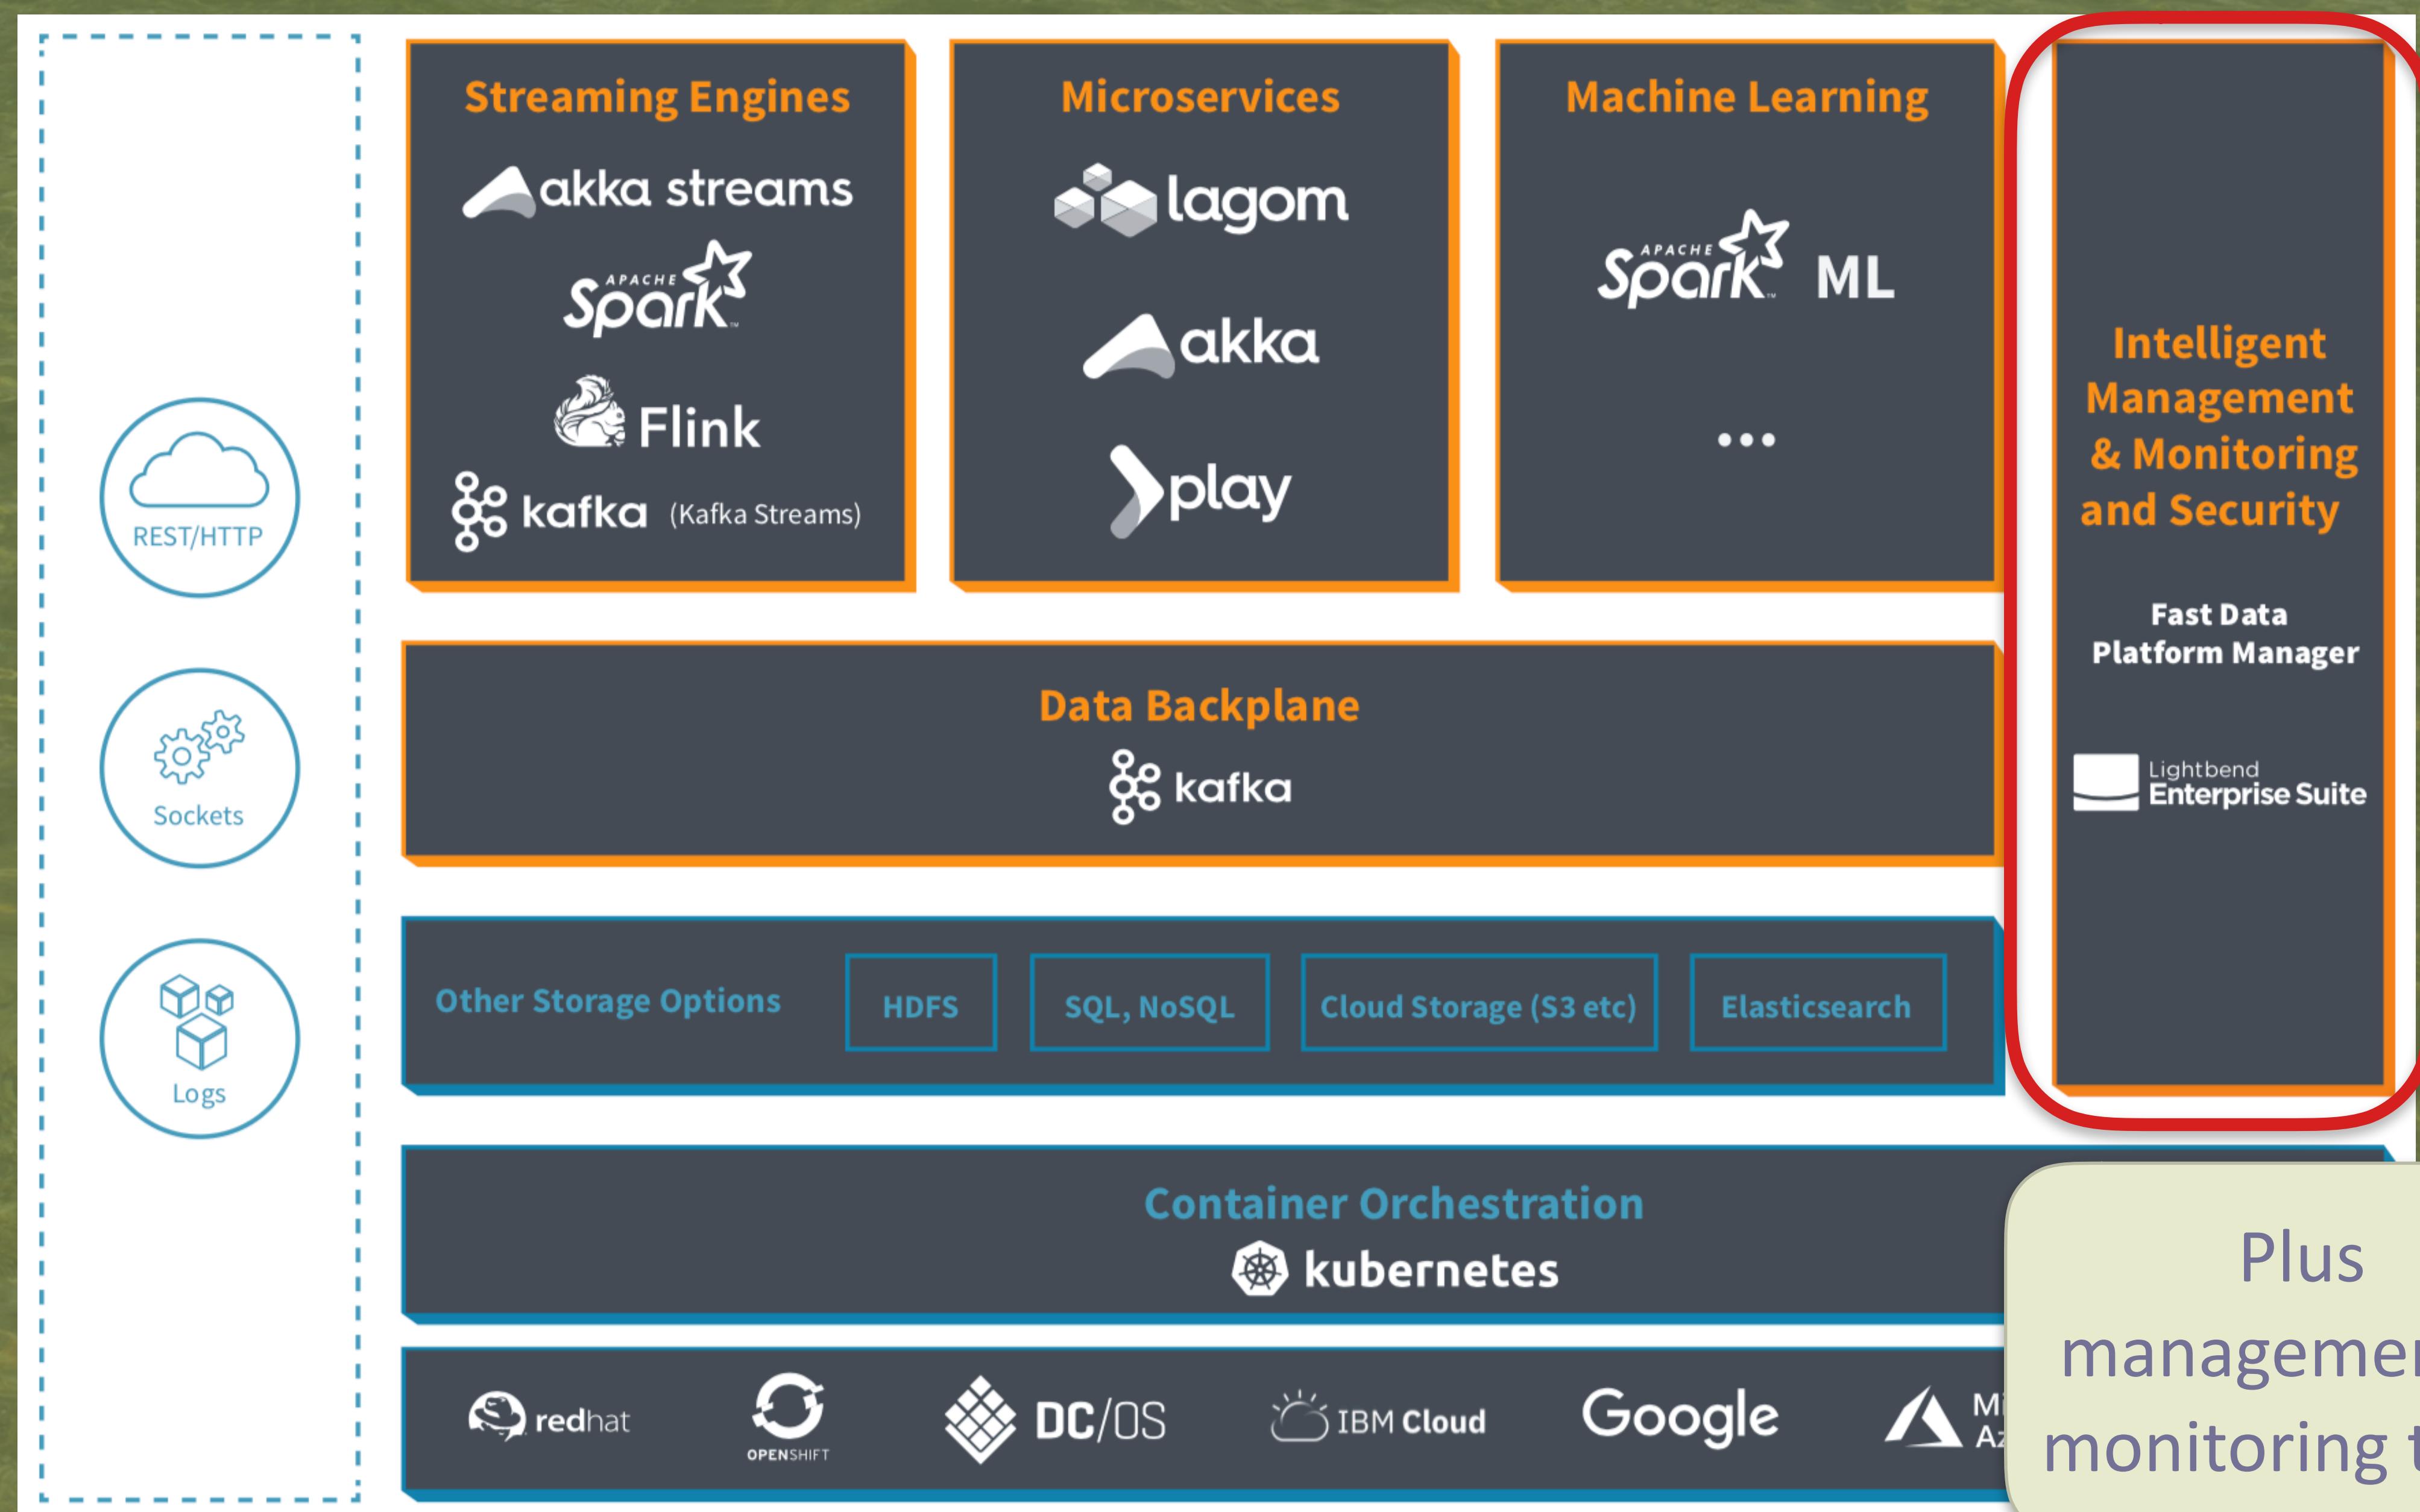


lightbend.com/fast-data-platform

What we
discussed



lightbend.com/fast-data-platform



lightbend.com/fast-data-platform



Lightbend

Questions?

Dean Wampler, Ph.D.

dean@lightbend.com

@deanwampler

polyglotprogramming.com/talks

lightbend.com/fast-data-platform